

claude-windows / c1fcb4aa-537d-41fd-858e-53f93349bf5a

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\subagents\workflows\wf_41c8dc8a-8a4\agent-a4bba0dd18652fded.jsonl`  
- SHA-256: `8f0766c0aaf51afc2fd357df80a04892fec0290c07b3804f741b94309a5c3764`  
- Source modified: `2026-07-02T03:03:48+00:00`  
- Imported at: `2026-07-05T16:48:24+00:00`  
- project: `wf_41c8dc8a-8a4`  
- session_id: `c1fcb4aa-537d-41fd-858e-53f93349bf5a`
```

Transcript

1. user (2026-07-02T03:01:38.608Z)

You are an adversarial verifier in a tech-debt audit. Your default stance: REFUTE. Today is 2026-07-02.

A finder claims the following about the repo at C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer:

```
{  
  "title": "deploy/digitalocean/ retired (ADR-013, droplet gone 2026-06-25) but still live-coupled: carries the trackers pin and an unverified 'rotate the key shared in chat' note",  
  "category": "doc-rot",  
  "severity": "P2",  
  "file": "deploy/digitalocean/README.md",  
  "line": 3,  
  "evidence": "deploy/digitalocean/README.md:3-9 marks DO DEPRECATED but says bootstrap.sh/mount_scratch.sh/gpu_setup.sh are 'still reused by the GCP runbook' and suggests moving them to deploy/common/; bootstrap.sh:46 is one of the 5 lockstep copies of the trackers pin; README.md:29 says 'Rotate the old key that was shared in chat' ? rotation is unverifiable from the repo (hypothesis: may never have happened).",  
  "why_it_matters": "A retired provider dir that live code paths depend on confuses the deploy story and keeps a possibly-unrotated-key instruction buried where nobody looks; ties into the pending post-droplet cleanup (api.khelsutra.guru CNAME/CF-Access) from project memory.",  
  "proposed_fix": "Move the three provider-agnostic scripts to deploy/common/ per the README's own suggestion, delete the DO-specific remainder, and have the owner confirm (or perform) the Gemini key rotation the README demands; also clears one copy of the trackers pin.",  
  "effort": "S",  
  "tracked_in": "deploy/digitalocean/README.md:8 (self-noted 'optional later cleanup')"  
}
```

Repo root candidates: C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer, C:/Users/avidu/Projects/khelsutra-guru/khelsutra, C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector, C:/Users/avidu/Projects/khelsutra-guru/rally-annotator, C:/Users/avidu/Projects/khelsutra-guru/sports-obsreport (pick from the finding's repo field; non-git dirs: 'C:/Users/avidu/Projects/khelsutra-guru/Annotation Setup', C:/Users/avidu/Projects/khelsutra-guru/khelsutra-screencast).

VERIFY against the CURRENT working tree (just pulled to origin master/main): (1) does the cited code actually exist and behave as claimed at file:line? (2) was it already fixed by a recent commit (git log --oneline -15 -- <file>)? (3) is the severity honest - construct the concrete failure scenario or downgrade; (4) is the proposed fix sane and correctly scoped; (5) is it owner-gated per: alpha/serving go-live, Studio P10/P11 corpus-promotion/training, product/strategy/licensing, real cloud spend, counsel ToS/DPA?

HARD RULES: READ-ONLY (no file modifications, no installs, no test-suite runs). gh read-only commands allowed. Your reasoning MUST cite file:line you actually read. CONFIRMED requires you personally saw the defect/debt in current code; if you cannot reproduce the claim from the code, say REFUTED; if a commit fixed it, ALREADY_FIXED.

2. user (2026-07-02T03:01:45.348Z)

```
1      # Testing the rally pipeline on DigitalOcean
2
3      > **? DEPRECATED (ADR-013, 2026-06-20) ? DigitalOcean is DROPPED for compute.** Use
4      > **[`deploy/gcp/README.md`](../gcp/README.md)** ? GCP is the sole compute stack (DO GPU
isn't enabled
5      > on the account + isn't cleanly credit-funded; Paperspace is subscription-gated). Kept
for history.
6      > **NOTE:** the scripts here ? `bootstrap.sh`, `mount_scratch.sh`, `gpu_setup.sh` ? are
7      > **provider-agnostic** Linux-GPU setup and are **still reused by the GCP runbook**;
```

only the

```
8      > DO-specific *provisioning* (doctl droplet create, Spaces) is retired. (Optional later
cleanup: move
9      > them to `deploy/common/`.)
10
11     A storage-free, secret-free way to get the pipeline running on a DigitalOcean (or
Paperspace)
12     machine ? the **compute** half of [`docs/archives/past_projects/DECOUPLED_COMPUTE/`](../
../docs/archives/past_projects/DECOUPLED_COMPUTE/README.md).
13     The object-store **storage shim is intentionally out of scope here** (owner is
finalizing it); this
14     runbook only proves *compute* works on the box.
15
16     > ? The commands below mirror `.github/workflows/ci.yml` (Ubuntu) so the box matches CI,
but the
17     > `bootstrap.sh` script has **not yet been run on a real DO box** ? treat it as a
tested-by-mirroring
18     > starting point and adjust as you go.
19
20     ## 1. Credentials to create (least-privilege)
21
22     | # | Credential | Where to create | Scope | Needed |
23     |---|---|---|---|---|
24     | 1 | **SSH keypair** | local `ssh-keygen`; add the **public** key to DO (Settings ?
Security) / Paperspace | machine login | **now** |
25     | 2 | **DigitalOcean API token** | DO console ? API ? Tokens ? Generate (Personal Access
Token) | **read + write** (Droplets/Spaces); scope down if the UI allows | **now** (for `doctl`;
optional if you click-create the box) |
26     | 3 | **Paperspace / Gradient API key** | console.paperspace.com ? API Keys | account |
**for the GPU box** (free M4000 / cheap A4000 live under Paperspace) |
27     | 4 | **GitHub repo read access** on the box | a **read-only deploy key** (repo ?
Settings ? Deploy keys) *or* a **fine-grained PAT** with `Contents: Read` | this repo only |
**now** (to `git clone`) |
28     | 5 | **DO Spaces keys** (access key + secret) | DO console ? API ? **Spaces Keys** |
one bucket/prefix | **defer** ? only when the storage shim + a real data test land |
29     | 6 | **`GEMINI_API_KEY`** | Google AI Studio (or Vertex) | the Gemini segmenter |
**later** ? only to test the *Gemini* path; the stub/WASB path needs none. **Rotate the old key
that was shared in chat.** |
30
31     **The "API key that drives the GPU" in the North-Star sense is **2 (DO token)** +
**3 (Paperspace key)** ?
32     they let you (or `doctl`/a control plane) spin up and drive GPU machines. The per-
invocation *serverless*
33     endpoint key (RunPod / Cloud Run) is a later M3 concern.
34
35     **Secret hygiene:** never bake any of these into the repo or a Docker image. Put them in
the box's
36     environment / a `.env` that is git-ignored, or the provider's secret store. Prefer
scoping #2 and #4 to the
37     narrowest resource the UI allows.
38
39     ## 2. Bring up a box
40
41     - **CPU smoke (M0/M1-cpu):** the cheapest **Droplet** (a basic shared-CPU Ubuntu
```

```

22.04/24.04) ? or, free, the
42     **Paperspace Gradient M4000** GPU machine (unbilled on the Free plan).
43     - **GPU test (M1 / M3.5 WSL?native parity):** a **Paperspace GPU** (free **M4000**, then
    **P4000/RTX4000/A4000**
44     on credit) or a **GPU Droplet**. All are Ubuntu + NVIDIA driver + CUDA.
45
46     `doctl` (optional) once #2 is set: `doctl auth init` then `doctl compute droplet create
?`. Or just create it
47     in the console and copy the IP.
48
49     ## 3. Bootstrap + smoke (on the box)
50
51     ```bash
52     ssh user@<box-ip>
53     git clone <repo>          # via deploy key (#4)
54     cd badminton-highlight-indexer
55
56     ./deploy/digitalocean/bootstrap.sh          # CPU box
57     # or
58     ./deploy/digitalocean/bootstrap.sh --gpu    # GPU box (installs CUDA torch)
59
60     . .venv/bin/activate && python -m pytest -q  # storage-free, secret-free smoke (~900
passed)
61     ```
62
63     **Acceptance:**
64     - CPU box ? the mocked suite passes (the whole pipeline runs with no
GPU/WSL/bucket/key).
65     - GPU box ? the bootstrap device check prints `cuda_available True` + the GPU name (M1
acceptance).
66
67     Once [#247](https://github.com/avidullu/badminton-highlight-indexer/pull/247) (the CPU-
mock detector seam)
68     is merged, you can also drive the trajectory segmenter GPU-free end-to-end:
69     `RALLY_DETECTOR_IMPL=stub python -m pytest tests/test_stub_runner.py -q`.
70
71     ## 4. Next compute steps (not in this runbook)
72
73     - **M3.5 WSL?native detector parity** ? replicate the `TRACKNET_WSL_SETUP.md` WASB conda
env *natively*
74     on the GPU box (no WSL), add a `LinuxNativeRunner`, and parity-test its trajectory CSV
vs the WSL box.
75     - **`train` / `infer` dispatcher + the bucket** ? gated on the storage shim (owner-
finalizing).
76

```

3. user (2026-07-02T03:01:46.221Z)

No files found

4. user (2026-07-02T03:01:51.931Z)

```

C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/deploy:
cloudrun
digitalocean
gcp

```

```

C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/deploy/digitalocean:
README.md
bootstrap.sh
gpu_setup.sh
mount_scratch.sh

```

5. user (2026-07-02T03:01:53.587Z)

af1ccbe chore(#399): launch-PR process hardening ? link-check CI, PR template, doc cross-links (#402)

3624a8d feat(D13/P12): migrate ByteTrack to standalone trackers package (#386)

2afbadd fix(dryrun): install the watermark font in bootstrap + track the UI-trigger asks (#292)

b05401b docs: drop DigitalOcean ? GCP is the sole compute stack (ADR-013) (#266)

0d73d96 docs(deploy): first real GCP-L4 run report + gpu_setup conda-ToS/deps fix + Ops Agent runbook (#261)

814d813 feat(deploy): GPU machine provisioning ? NVMe scratch + native WASB setup (#253)

bf9776f chore(deploy): DigitalOcean compute bootstrap + runbook (storage-free) (#248)

e410136 api: /api/videos surfaces fps + resolution (from the stored ffprobe details) (#413)

81a78bf api: /api/videos returns video_id + a friendly match_name (fix disappearing name) (#412)

8b16cf4 docs: reflect today's storage endpoints + the corpus/user-video boundary (#411)

6. user (2026-07-02T03:02:02.062Z)

```

1      #!/usr/bin/env bash
2      #
3      # Bootstrap a DigitalOcean / Paperspace Ubuntu box to run the rally pipeline.
4      #
5      # ./deploy/digitalocean/bootstrap.sh          # CPU box (Droplet)          ? M0 /
M1-cpu
6      # ./deploy/digitalocean/bootstrap.sh --gpu    # GPU box (Paperspace / GPU Droplet) ?
M1 / M3.5
7      #
8      # Mirrors .github/workflows/ci.yml (Ubuntu) so the box matches CI exactly.
9      # STORAGE-FREE + SECRET-FREE: proves the pipeline runs on the box via the mocked
10     # test suite. No bucket, no Gemini key, no WSL required. See
11     ../../docs/DECOUPLED_COMPUTE/.
12     set -euo pipefail
13     GPU=0
14     [ "${1:-}" = "--gpu" ] && GPU=1
15
16     echo ">> [1/4] System libraries (ffmpeg + opencv runtime + the watermark font)"
17     sudo apt-get update
18     # fonts-dejavu-core: the stitcher's watermark uses ffmpeg drawtext with DejaVu ? WITHOUT
it
19     # compile_highlights raises and NO reel is produced (found on the 2026-06-21 cloud dry
run: 22
20     # rallies detected but no reel). Matches deploy/cloudrun/Dockerfile. (Or set
stitching.watermark=false.)
21     sudo apt-get install -y python3 python3-venv python3-pip git ffmpeg libgl1 fonts-dejavu-
core
22     # opencv needs libgl2.0-0; the package name differs across Ubuntu releases (t64 on
24.04).
23     sudo apt-get install -y libgl2.0-0t64 || sudo apt-get install -y libgl2.0-0
24
25     echo ">> [2/4] Python venv (.venv)"
26     python3 -m venv .venv
27     # shellcheck disable=SC1091
28     . .venv/bin/activate
29     python -m pip install --upgrade pip
30
31     echo ">> [3/4] torch (${[ "$GPU" = 1 ] && echo CUDA || echo CPU}) + package (editable,
dev tooling)"
32     if [ "$GPU" = 1 ]; then
33         # CUDA 12.4 wheels. NOTE: verify your GPU arch is supported by this build.
34         # The Paperspace free M4000 is Maxwell (sm_52) ? old archs are sometimes dropped;
35         # if `torch.cuda.is_available()` is False on an M4000, use a P4000/RTX4000 or pin
36         # an older torch. (docs/DECOUPLED_COMPUTE/03 + 06.)
37         pip install torch torchvision --index-url https://download.pytorch.org/whl/cu124
38     else
39         pip install torch torchvision --index-url https://download.pytorch.org/whl/cpu
40     fi
41     # D13/P12: ByteTrack moved out of `supervision` (removed in 0.30.0) to the standalone

```

```

42      # `trackers` package. Git-only dep (PyPI wheel is metadata-only ?
roboflow/trackers#474),
43      # so it's not in pyproject.toml; install out-of-band exactly as ci.yml/nightly.yml do,
or
44      # the person-cue path (yolo_hybrid/fusion_hybrid) crashes at the first chunk with
45      # ModuleNotFoundError: No module named 'trackers'. Bare version tag (roboflow uses no
'v').
46      pip install "trackers @ git+https://github.com/roboflow/trackers.git@2.5.0"
47      # Runtime deps + pytest/ruff/mypy. The 'reporting' extra (private obsreport over SSH)
48      # is intentionally omitted ? it is a soft dep and its tests skip when absent.
49      pip install -e .[dev]
50
51      echo ">> [4/4] Device check"
52      python - <<'PY'
53      import torch
54      ok = torch.cuda.is_available()
55      print("torch", torch.__version__, "| cuda_available", ok,
56            "| device", (torch.cuda.get_device_name(0) if ok else "CPU only"))
57      PY
58
59      cat <<'EOF'
60
61      >> Bootstrap complete. Run the storage-free, secret-free smoke (mocked suite):
62
63          . .venv/bin/activate && python -m pytest -q
64
65          Expect ~900 passed. That proves the whole pipeline runs on this box with no GPU,
66          no WSL, no bucket and no API key. (On a GPU box the device check above should
67          print cuda_available True.)
68      EOF
69

```

7. user (2026-07-02T03:02:03.426Z)

```

badminton-highlight-indexer\deploy\gcp\README.md:123:> `nvidia-smi` works immediately) ? used
for the first real L4 run. `gpu_setup.sh` installs miniconda
badminton-highlight-indexer\deploy\gcp\README.md:179:sudo bash
deploy/digitalocean/mount_scratch.sh && export TMPDIR=/scratch # detects the local NVMe
badminton-highlight-indexer\deploy\gcp\README.md:180:./deploy/digitalocean/bootstrap.sh --gpu
# venv + torch cu124 (confirm cuda True)
badminton-highlight-indexer\deploy\gcp\README.md:182:bash deploy/digitalocean/gpu_setup.sh
# native `wasb` conda env + WASB-SBDT
badminton-highlight-indexer\deploy\gcp\README.md:191:~/models/WASB-SBDT/pretrained_weights/`
(gpu_setup.sh does NOT fetch them ? stage from Drive/bucket).
badminton-highlight-indexer\deploy\gcp\nightly_regression\run_on_vm.sh:86:sudo bash
deploy/digitalocean/mount_scratch.sh 2>/dev/null || true
badminton-highlight-
indexer\deploy\gcp\nightly_regression\run_on_vm.sh:88:./deploy/digitalocean/bootstrap.sh --gpu
badminton-highlight-indexer\deploy\gcp\nightly_regression\run_on_vm.sh:91:bash
deploy/digitalocean/gpu_setup.sh || echo "WARN: gpu_setup (WASB env) returned nonzero ?
continuing"
badminton-highlight-indexer\deploy\gcp\nightly_regression\run_on_vm.sh:93:# 4) Stage the WASB
checkpoint (gpu_setup does NOT fetch weights ? README §4) + export the FULL native-
badminton-highlight-indexer\deploy\gcp\nightly_regression\run_on_vm.sh:96:# L4, 2026-06-22). The
conda env is built by gpu_setup.sh under $HOME/miniconda3/envs/wasb.

```

8. assistant (2026-07-02T03:02:13.922Z)

The core claims check out so far. Now let me verify the "5 lockstep copies of the trackers pin" claim and the ADR-013 / droplet-gone dates.

9. user (2026-07-02T03:02:15.477Z)

```

badminton-highlight-indexer\requirements.txt:22:# pip install "trackers @

```

```

git+https://github.com/roboflow/trackers.git@2.5.0"
badminton-highlight-indexer\requirements.txt:29:trackers @
git+https://github.com/roboflow/trackers.git@2.5.0
badminton-highlight-indexer\.github\workflows\ci.yml:132:      run: python -m pip install
"trackers @ git+https://github.com/roboflow/trackers.git@2.5.0"
badminton-highlight-indexer\.github\workflows\nightly.yml:135:      run: python -m pip install
"trackers @ git+https://github.com/roboflow/trackers.git@2.5.0"
badminton-highlight-indexer\deploy\cloudrun\Dockerfile:62:      && python3 -m pip install
"trackers @ git+https://github.com/roboflow/trackers.git@2.5.0" \
badminton-highlight-indexer\docs\CODE_CLEANUP_AUDIT_2026-06-24.md:282:[Omitted long matching
line]
badminton-highlight-indexer\deploy\digitalocean\bootstrap.sh:46: pip install "trackers @
git+https://github.com/roboflow/trackers.git@2.5.0"

```

10. user (2026-07-02T03:02:16.892Z)

```

badminton-highlight-indexer\docs\DOC_STATUS.md:61:| `HOSTING_PLAN.md` | PROPOSAL | **? partly
overtaken by ADR-013 + the shipped site ? verify (finding #5)** |
badminton-highlight-indexer\docs\DOC_STATUS.md:91:5. **`HOSTING_PLAN.md` / `UI_PROPOSAL.md` are
PROPOSAL but partly overtaken** by ADR-013 (DO dropped) + the shipped khelsutra site. **?
`UI_PROPOSAL.md` banner added (2026-06-21)** ? its UX half is now tracked in
`khelsutra/docs/FRIENDLY_AND_MULTILINGUAL.md`; the `I18N_PLAN.md` localization track is likewise
in implementation there. **`HOSTING_PLAN.md` banner still pending** (ADR-013 already noted in
it).
badminton-highlight-indexer\docs\HOSTING_PLAN.md:7:> **? Update (2026-06-20, ADR-013):
DigitalOcean is DROPPED for COMPUTE** ? GCP is the sole compute
badminton-highlight-indexer\docs\HOSTING_PLAN.md:78:| **All-DigitalOcean** | ~$11 ? ~$17 ($6
droplet + Spaces $5: 250 GiB + **1 TiB egress included**) | One simple vendor; Spaces' included
egress is the best video-download deal; per-second billing | No credits; no scale-to-zero; still
needs a tunnel-equivalent to reach the box (run `frp`/WireGuard on the droplet) |
badminton-highlight-indexer\docs\COMPUTE_DECOUPLED_SERVING\architecture.svg:128: <text x="32"
y="688" font-size="11" fill="#475569"><tspan font-weight="700">Lineage:</tspan> ADR-009 (local-
first) ? ADR-010/011 (decoupled; M3/M4 deferred) ? ADR-012 (Cloud Run+GPU) ? ADR-013 (GCP sole
stack) ? ADR-014 (this).</text>
badminton-highlight-indexer\docs\COMPUTE_DECOUPLED_SERVING\README.md:57:| **ADR-013** | Drop DO;
**GCP sole compute stack** ; $200 DO credits ? non-compute only | The provider for cloud-serving;
DO Spaces remains only a possible S3 *storage* target. |
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:528:- **DigitalOcean GPU
Droplets are North-America-only** (`nyc2`/`tor1`/`atl1`) ? DO has **no
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:530: droplet-in-`blr1` +
Spaces-in-`sgp1` ? cross-DC latency + egress. GCP has GPUs in asia-south1
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:536:- The old **DO CPU droplet
(id `579001481`, `blr1`) is drained** (`/srv/khelsutra-site` copied off)
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:545:- `SETUP_NEW_MACHINE.md`
$2's "DO GPU droplet in `sgp1`" is **superseded** by the GCP path
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:553:**Date:** 2026-06-20 .
**Status:** RATIFIED (owner, in-session). **Refines ADR-010.** ? **The DO=training / GCP=serving
split is SUPERSEDED by [ADR-013](#adr-013-drop-digitalocean-for-compute--gcp-is-the-sole-
compute-stack-training--serving)** ? DigitalOcean is dropped for compute; GCP is the sole
compute stack. (The scope decision ? DECOUPLED = M0?M2 + M3.5 ? still stands.)
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:566: - **DigitalOcean =
TRAINING / the decoupling core.** Use the **$200 DO credits** ; DO GPU droplets
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:595: staged to Spaces. The old
DO CPU droplet `579001481` is drained + destroyed; the exposed Spaces key
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:661:## ADR-013: Drop
DigitalOcean for compute ? GCP is the SOLE compute stack (training + serving)
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:669:3. The `$200` DO credits,
if used at all, go **only to non-compute** that DO bills cleanly (CPU droplets / a site host /
Spaces storage / bandwidth) ? **never GPU** .
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:673:[Omitted long matching
line]
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:674:- **Paperspace (now DO
Gradient) free GPU is subscription-gated** (a paid "G"/Growth plan; high-end A100/H100 needs
the $39/mo Growth plan), and Paperspace bills **separately** from DO Droplet credits ? the $200
credit doesn't fund it either.

```

badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:680:- ****`deploy/digitalocean/`** scripts are KEPT**** ? `bootstrap.sh` / `mount\_scratch.sh` / `gpu\_setup.sh` are ****provider-agnostic**** Linux-GPU setup the GCP runbook reuses; only the DO-specific ***provisioning*** (doctl droplet create, Spaces) is dropped. (Optional later cleanup: move them to `deploy/common/`.)**

badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:681:- The shared DO CPU droplet ****`claude-do-rally-test`**** (another session's testbed) is ****owner-coordinated**** cleanup, separate from this decision; the freshly-rotated ****DO API token**** can be retired once that droplet is gone.

badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:688:****Date:** 2026-06-21 .**
****Status:** ? ****RATIFIED** (owner, 2026-06-21 ? M0 design sign-off; build M1+). ****Refines** ADR-011 + ADR-012****** (it ***realizes*** their deferred M3/M4; it does ****not**** supersede them); ****builds on ADR-013**** (GCP sole compute stack). Full design: `docs/COMPUTE\_DECOUPLED\_SERVING/`.**

badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:691:[Omitted long matching line]

badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:707:- Aligns with ADR-012 (GPU-native) + ADR-013 (GCP sole stack) + the 2026-06-20 deploy report; the truly-local profile preserves the ADR-009 privacy/offline self-host story.

badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:75:| D3 | Runs on an ****ephemeral** GCP GPU VM that self-deletes****** (not GitHub CI, not local) | owner 2026-06-22 + ADR-013 (compute=GCP) | GitHub CI can't run WASB (no GPU/weights); the VM is the only faithful venue; self-delete = cost rule |

badminton-highlight-indexer\docs\NEXT_STEPS.md:6:[Omitted long matching line]

badminton-highlight-indexer\docs\NEXT_STEPS.md:33:15. Owner housekeeping: retire the shared ****DO droplet**** + ****DO token****; collector ****`md5`**** export (collector repo).

badminton-highlight-indexer\docs\archives\past_projects\README.md:12:- **`DECOUPLED\_COMPUTE/` ? ? ****DELIVERED + ARCHIVED 2026-06-21.**** Lifted the rally pipeline onto a rented cloud GPU (bucket-driven train + infer); ****M0?M2 + M3.5 proven on a GCP L4**** (M2 ? `weights/0.1.0-l4/`). Delivered on ****GCP**** ? DigitalOcean dropped for compute (ADR-013). ****M3/M4** (serving endpoint + per-video billing + scale-to-zero) deferred ? the "Cloud Run + GPU serving" subproject****** (ADR-012). The full deploy record is **`DECOUPLED\_COMPUTE/DEPLOY\_REPORT\_GCP\_2026-06-20.md`**.**

badminton-highlight-indexer\docs\archives\past_projects\PACKAGING_PLAN.md:38:| L3 | ****Compute = GCP only**** (Cloud Run GPU Jobs); DigitalOcean dropped | ADR-013, 2026-06-20 | Cloud tier targets Cloud Run + BYO-GPU-VM only. |

badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\03-PLATFORM-OPTIONS.md:16:| ****Plain VM + `venv/`/`conda`**** | A DO GPU Droplet or GCE VM you **`git clone`** + **`pip install -r requirements-train.txt`** onto. Simplest first step; matches today's dev flow. | No image to build, but ****you**** manage CUDA drivers + reproducibility is weaker. Fine for M1?M2. |

badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\03-PLATFORM-OPTIONS.md:17:| ****SkyPilot**** (already named in [PLATFORM_ARCHITECTURE.md:156-160])(../PLATFORM_ARCHITECTURE.md)) | "No-Docker" managed provisioning: a YAML spec launches your repo in a conda env on a spot VM on a backing cloud, handles bucket mounts + auto-stop. | Backs GCP/AWS/Azure/Lambda/RunPod/Paperspace (****not**** DO Droplets natively). Great cost-control with zero image work. |

badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\03-PLATFORM-OPTIONS.md:40:[Omitted long matching line]

badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\03-PLATFORM-OPTIONS.md:55: ****Paperspace/Gradient A4000**** (or a small GPU Droplet) running the repo in a ****venv**** (no Docker needed

badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\03-PLATFORM-OPTIONS.md:69:a ****DO GPU Droplet**** or ****Paperspace/Gradient GPU**** is ****Ubuntu + NVIDIA + CUDA****, so you replicate the

badminton-highlight-indexer\docs\README.md:18:[Omitted long matching line]

badminton-highlight-indexer\docs\README.md:25:[Omitted long matching line]

badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\04-BUILD-AND-TEST-PLAN.md:111:AI handoff, eval) run ****verbatim, byte-for-byte****. On a ****DO GPU Droplet / Paperspace GPU**** (Ubuntu +

badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\05-COST-ATTRIBUTION.md:30:| ****Worker-measured `gpu\_active\_seconds`**** (timestamp around the GPU section, × the known **`gpu\_type`** rate) | DO Droplet / VM where there's no serverless meter | Good ? separate **`gpu\_active\_seconds`** from **`wall\_seconds`** so you don't bill the user for the bucket-sync/upload time. |

badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\05-COST-ATTRIBUTION.md:42:ALTER TABLE jobs ADD COLUMN compute_backend TEXT; -- runpod_serverless | gcp_cloudrun | do_droplet ?

```
badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\07-COMPUTE-ABSTRACTION.md:31: per-job: local | DO droplet ? per-stage: probe the actual GPU ?
badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\07-COMPUTE-ABSTRACTION.md:38:this stage*. ** A serverless L4, a DO RTX-4090 Droplet, and CPU CI each get correct knobs **automatically**,
badminton-highlight-indexer\docs\SETUP_NEW_MACHINE.md:6:> **? COMPUTE = GCP (ADR-013, 2026-06-20). DigitalOcean is DROPPED for compute** ? use
```

[Showing results with pagination = limit: 40]

11. user (2026-07-02T03:02:26.394Z)

```
badminton-highlight-indexer\deploy\digitalocean\README.md:29:| 6 | **`GEMINI_API_KEY`** | Google AI Studio (or Vertex) | the Gemini segmenter | **later** ? only to test the *Gemini* path; the stub/WASB path needs none. **Rotate the old key that was shared in chat.** |
badminton-highlight-indexer\deploy\gcp\README.md:184:# secrets (NOT in repo): the GCS SA JSON + a rotated Gemini key
badminton-highlight-indexer\deploy\gcp\README.md:187:mkdir -p /root/.keys && printf '%s' '<ROTATED_GEMINI_KEY>' > /root/.keys/GEMINI_API_KEY && chmod 600 /root/.keys/GEMINI_API_KEY
badminton-highlight-indexer\deploy\gcp\provision.sh:56: echo " key already at $SA_KEY ? leaving it (delete to rotate)"
badminton-highlight-indexer\docs\COMPUTE_DECOUPLED_SERVING\README.md:170:[Omitted long matching line]
badminton-highlight-indexer\backend\api\auth.py:26: 600.0 # re-fetch the team JWKS at most every 10 min (keys rotate slowly)
badminton-highlight-indexer\docs\GOLDEN_REGRESSION_FIXTURES.md:113:[Omitted long matching line]
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:537: **+ destroyed**; the **exposed DO Spaces key is rotated** (it leaked in chat).
badminton-highlight-indexer\docs\archives\decisions\DECISIONS.md:594:- The DO training run needs: a DO API token (owner), the rotated Gemini key, WASB weights, corpus
badminton-highlight-indexer\docs\DATA_IN_GCS.md:228: created ? `deploy/gcp/README` $2); cloud-box secrets (SA JSON + rotated `GEMINI_API_KEY` + WASB
badminton-highlight-indexer\docs\HOSTING_PLAN.md:110:allowlist ? **rotate the Gemini API key** (NEXT_STEPS housekeeping ? it was shared in chat) ?
badminton-highlight-indexer\docs\NEXT_STEPS.md:149:6. **Housekeeping:** ~rotate the Gemini API key~ **DONE 2026-06-10** (owner-verified working; confirm old-key
badminton-highlight-indexer\docs\UI_ALPHA_EXECUTION_PLAN.md:25:- [x] **Owner:** Gemini API key rotated ? 2026-06-10, owner-verified (confirm old-key deletion in AI Studio)
badminton-highlight-indexer\docs\SETUP_NEW_MACHINE.md:36:| 4 | **`GEMINI_API_KEY`** | Google AI Studio / Vertex | **real** Gemini inference only ($6b). The secret-free path ($6a) needs none. **Rotate before sharing.** |
badminton-highlight-indexer\docs\SETUP_NEW_MACHINE.md:162:export GEMINI_API_KEY="<your-rotated-key>" # secret stays in the env, never in the repo
badminton-highlight-indexer\docs\SETUP_NEW_MACHINE.md:272:mkdir -p /root/.keys && printf '%s' '<rotated-key>' > /root/.keys/GEMINI_API_KEY && chmod 600 /root/.keys/GEMINI_API_KEY
badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\HANDOFF.md:23: drained (`/srv/khelsutra-site` copied off) + destroyed; the exposed DO Spaces key is rotated.
badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\HANDOFF.md:69: ? **The Spaces key was exposed in chat ? ROTATE it.** Creds live on the box at `~/aws/credentials` (never in repo).
badminton-highlight-indexer\docs\archives\past_projects\DECOUPLED_COMPUTE\HANDOFF.md:88: (native WASB conda env) ? secrets (rotated Gemini key + Spaces creds).
badminton-highlight-indexer\docs\UI_PROPOSAL.md:66:| B7 | **Exposure hardening** | CORS narrowed to the app origin; basic rate limit; upload quota per user; `allowed_video_root` set; Gemini key rotated (already in NEXT_STEPS housekeeping). |
badminton-highlight-indexer\backend\ui\assets\index-Jp9DK5o0.css:1:[Omitted long matching line]
badminton-highlight-indexer\tests\test_run_telemetry.py:115:def test_rotation_caps_size_keeps_header_and_recent(tmp_path, monkeypatch):
```

12. assistant (2026-07-02T03:03:06.090Z)

Let me read the two rotation-status lines directly to confirm context.

13. user (2026-07-02T03:03:07.508Z)

145 `docs/FIELD\_VISIT\_PLAYBOOK.md` (the hired associate's visit playbook: coach + student
script, demo-after-Q6
146 rule, ground rules, Phase-2 record-and-show-back stretch goal) +
`docs/FIELD\_VISIT\_ANSWER\_SHEET.md` (per-visit
147 capture form incl. demo-reaction + student blocks) + `docs/PMF\_FIELD\_SIGNALS.md`
(G/A/R anchors + **pre-registered
148 decision rules** ? PROCEED / PIVOT A-C / FALSIFIED thresholds fixed before the first
visit).
149 6. **Housekeeping:** ~rotate the Gemini API key~~ **DONE 2026-06-10** (owner-verified
working; confirm old-key
150 deletion in AI Studio); **rename repo off "badminton"** (owner action,
151 in BACKLOG); collector=SoR confirmation; off-peak niceties: none pending ? the
battery completed.
152

14. user (2026-07-02T03:03:08.231Z)

22 - [x] **Owner:** git-init + private GitHub push ? done 2026-06-10:
23 `https://github.com/avidullu/khelsutra` (private; git runs Windows-side per ADR-009)
24 - [x] **Owner:** explicit "go" for PR-1 ? given 2026-06-10
25 - [x] **Owner:** Gemini API key rotated ? 2026-06-10, owner-verified (confirm old-key
deletion in AI Studio)
26
27 ### Phase L ? build & test locally (no cloud; ?6?7 sessions)

15. user (2026-07-02T03:03:48.630Z)

Structured output provided successfully