

claude-windows / c1fcb4aa-537d-41fd-858e-53f93349bf5a

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\subagents\workflows\wf_41c8dc8a-8a4\agent-a8a3025ad2312e3de.jsonl`  
- SHA-256: `d7675576f3098c28af147e84de92e5e934cc3f5522e64c0939cf5704b42b320b`  
- Source modified: `2026-07-02T03:01:07+00:00`  
- Imported at: `2026-07-05T16:48:24+00:00`  
- project: `wf_41c8dc8a-8a4`  
- session_id: `c1fcb4aa-537d-41fd-858e-53f93349bf5a`
```

Transcript

1. user (2026-07-02T02:59:39.461Z)

You are an adversarial verifier in a tech-debt audit. Your default stance: REFUTE. Today is 2026-07-02.

A finder claims the following about the repo at C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer:

```
{  
  "title": "P0-rename remainder: person/venue-attributable video names still committed in code and baselines",  
  "category": "security",  
  "severity": "P2",  
  "file": "backend/eval/rally_seg_eval.py",  
  "line": 53,  
  "evidence": "backend/eval/rally_seg_eval.py:52-54: MULTICOURT_CLIPS =  
frozenset({'mahadevpura_2\\', '\\GX010128\\', '\\Badminton_BXH_2\\', '\\Boxhill_Doubles\\'}).  
eval_baselines/heuristic_lovo_n6.json per_video keys: adarsh_avi_singles, kushagra_singles,  
mahadevpura_singles, etc. (person first names + venues). Matches  
GOLDEN_REGRESSION_FIXTURES.md:151 P0 row (functional rename DEFERRED per owner safe-scope  
2026-06-16) and BACKLOG.md:61 ('Revisit: before any open-sourcing / public commit').",  
  "why_it_matters": "The repo's own leak-scan/de-attribution posture (fx_* surrogate slugs, CI  
leak-scan job) is undermined by these functional names: they attribute real people/venues in a  
repo the owner intends to rename and potentially share. The docs' claim is accurate ? this is a  
verified still-open gate on any public commit, and it silently couples to the gitignored LOVO  
manifest (rename breaks the training floor if done uncoordinated).",  
  "proposed_fix": "Before any open-sourcing or the repo rename: put the names in gitignored  
.leakscan_private, run python -m tools.leak_scan --names .leakscan_private to enumerate  
occurrences, then execute the coordinated rename (baseline keys + MULTICOURT_CLIPS + quality-log  
names + owner-side output/human_lovo_manifest.json together).",  
  "effort": "M",  
  "tracked_in": "docs/GOLDEN_REGRESSION_FIXTURES.md §7 P0; docs/BACKLOG.md:61"  
}
```

Repo root candidates: C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer,
C:/Users/avidu/Projects/khelsutra-guru/khelsutra, C:/Users/avidu/Projects/khelsutra-guru/sports-
data-collector, C:/Users/avidu/Projects/khelsutra-guru/rally-annotator,
C:/Users/avidu/Projects/khelsutra-guru/sports-obsreport (pick from the finding's repo field;
non-git dirs: 'C:/Users/avidu/Projects/khelsutra-guru/Annotation Setup',
C:/Users/avidu/Projects/khelsutra-guru/khelsutra-screencast).

VERIFY against the CURRENT working tree (just pulled to origin master/main): (1) does the cited
code actually exist and behave as claimed at file:line? (2) was it already fixed by a recent
commit (git log --oneline -15 -- <file>)? (3) is the severity honest - construct the concrete
failure scenario or downgrade; (4) is the proposed fix sane and correctly scoped; (5) is it
owner-gated per: alpha/serving go-live, Studio P10/P11 corpus-promotion/training,
product/strategy/licensing, real cloud spend, counsel ToS/DPA?

HARD RULES: READ-ONLY (no file modifications, no installs, no test-suite runs). gh read-only

commands allowed. Your reasoning MUST cite file:line you actually read. CONFIRMED requires you personally saw the defect/debt in current code; if you cannot reproduce the claim from the code, say REFUTED; if a commit fixed it, ALREADY_FIXED.

2. user (2026-07-02T02:59:50.115Z)

```
40     #: strips most of those FPS and reveals the real product quality on the long-rally happy
path. Used
41     #: as the default split point for the stratified report (`--long-tau` overrides it).
DURATION, not
42     #: shot-count, is the filter: golden `shots_count` is unlabeled and the no-AI path
reports it as
43     #: "Unknown", whereas duration is derived from the same start/end labels we already
trust.
44     LONG_RALLY_TAU = 5.0
45
46     #: Multicourt golden clips ? the court-type lookup for the stratified report. The golden
manifest
47     #: (`output/human_lovo_manifest.json`) is gitignored, so this committed set is the
AUDITABLE source
48     #: of the single-vs-multicourt strata, kept in sync with `docs/GOLDEN_VIDEOS.md` (the
clips tagged
49     #: "multicourt"). Multicourt footage carries background games on adjacent courts ? the
dominant
50     #: source of short false positives ? so we report it separately. A manifest entry MAY
carry an
51     #: explicit `court_type` field, which overrides this set (see :func:`court_type_for`).
52     MULTICOURT_CLIPS = frozenset(
53         {"mahadevpura_2", "GX010128", "Badminton_BXH_2", "Boxhill_Doubles"}
54     )
55
56
57     def court_type_for(video: Dict[str, Any]) -> str:
58         """`multicourt` or `single` for a golden video. Honors an explicit manifest
59         `court_type` field when present, else falls back to membership in
`MULTICOURT_CLIPS`
60         (keyed by `name`), else `single`."""
61         ct = video.get("court_type")
62         if isinstance(ct, str) and ct:
63             return ct
64         return "multicourt" if video.get("name") in MULTICOURT_CLIPS else "single"
65
66
67     def filter_by_min_duration(ivs: List[Interval], min_duration: float) -> List[Interval]:
68         """The long-rally lens: keep only intervals strictly longer than `min_duration`
seconds.
69
```

3. user (2026-07-02T02:59:51.611Z)

```
0c9c693 feat(eval): expose inactivity_temporal_density in the rally_seg_eval A/B CLI (P2a)
(#395)
b6214e5 chore: ruff format the whole repo (259 files) (#387)
e358b28 feat(eval): prominent-court window gate ? drop multicourt background-court rallies
(default-OFF) (#243)
20f8aa3 feat(eval): P7 long-rally (>5s) quality lens + reel knob (both default-OFF) (#220)
ae10447 refactor(eval): extract _build_preset(args, side) from rally_seg_eval.main() (#231)
f5f5e52 feat(eval): R5 ? blob-fallback splits the continuously-active blob (default-OFF) (#203)
8c68143 feat(eval): R1 ? net over-seg promotion guard unblocks the cap6+R4 milestone (#201)
34b9caa feat(eval): R4 ? rally-END inactivity trim (default-OFF) + R3 cap re-scope finding
(#200)
7e8f88f feat(eval): R2 ? task-faithful tolerance-match rally metric (augments tIoU) (#199)
bb4d6e2 feat(eval): Tier 1 (W1.5) ? hard-cap fallback enforces the window cap on continuously-
active trajectories (#195)
950106b fix(fusion): W2 net-crossings gate is do-no-harm ? never empty a non-empty video (#188)
```

```
follow-up) (#190)
a471745 feat(eval): Tier 1 (W2) ? net-crossings rally-state gate on W1 (?F1 +0.019, 6/0) (#188)
90ab087 feat(eval): Tier 0 (E1) ? over-segmentation-aware rally-seg eval + frozen baseline
(#185)
---
e410136 api: /api/videos surfaces fps + resolution (from the stored ffprobe details) (#413)
81a78bf api: /api/videos returns video_id + a friendly match_name (fix disappearing name) (#412)
8b16cf4 docs: reflect today's storage endpoints + the corpus/user-video boundary (#411)
aa915a7 deploy/gcp: centralize project + bucket names in one sourced buckets.env (#410)
52d97b5 P4: POST /api/assets ? store a video into the user's chosen medium (#409)
```

4. user (2026-07-02T02:59:58.206Z)

```
['adarsh_avi_singles', 'kushagra_singles', 'largetest_doubles', 'gbaaddy', 'testlarge_short',
'mahadevpura_singles']
```

5. user (2026-07-02T03:00:00.108Z)

```
21:- ? **Live de-attribution bug** *(verified)*: real player/venue names are committed today in
`eval_baselines/heuristic_lovo_n6.json`, `eval_baselines/gemini_wrapper_2026-06-10.json`, and
`backend/eval/eval_partial_labels.py` (the last also has a personal `C:/Users/avidu/...` path) ?
and in git history. ? **Deliverable P0.**
151:[Omitted long matching line]
164:[Omitted long matching line]
173:| _this PR (P0 safe scope)_ | personal-path scrub + leak-scan guard | 1365 pass / 7 skip |
85.27% | scrubbed `C:\Users\?` paths from 3 code docstrings (`eval_partial_labels.py` /
`gemini_refine.py` / `README.md`); `tools/leak_scan.py` (32-hex MD5 + personal-path structural
guard + gitignored `.leakscan_private` name-list) + a CI `leak-scan` job + 10 tests. Functional
video-name rename DEFERRED (owner safe-scope). |
197:- [ ] **P0, M1?M4, P1, DOC** are ? and merged.
```

6. user (2026-07-02T03:00:07.256Z)

```
140      - De-attribution does **not** cure provenance/licensing ? it sits *behind* the per-
record provenance gate + counsel sign-off.
141      - `--refresh-snapshot` can rubber-stamp a real regression; mitigation is **review of the
visible diff** + a separate, deliberate floor-baseline update. **A recorded `--reason` is
REQUIRED** (the CLI errors without one ? audit fix #6).
142
143      ---
144
145      ## 7. Deliverables & progress tracker ? **source of truth**
146
147      Legend: ? Todo . ? In progress . ? Done . ? Blocked/gated. One **small PR per row**,
branched from `origin/master`, no stacking. Update the row (status + PR link) as work lands.
148
149      | ID | Deliverable | Depends on | Counsel-gated? | Status | PR |
150      |----|-----|-----|-----|-----|----|
151      | **P0** | **PII leak hygiene (partial ? paths + scan).** Scrubbed personal `C:\Users\?`
paths from code docstrings (`eval_partial_labels.py`, `gemini_refine.py`, `README.md`) + added
`tools/leak_scan.py` (32-hex MD5 + personal-path **structural** guard, optional gitignored
`.leakscan_private` name-list) + a CI `leak-scan` job + tests. ? The **functional video-name
rename** (`eval_baselines/*.json` keys / `rally_seg_eval.MULTICOURT_CLIPS` / quality-log names)
is **DEFERRED** per the owner's safe-scope decision (2026-06-16) ? it's coupled to the
gitignored `output/human_lovo_manifest.json` + data files + the training floor and collides with
the active quality track; run the private-list scan + a coordinated rename before open-sourcing.
Git-history rewrite (PR #77) also deferred. | ? | No | ? | _this PR_ |
152      | **M1** | **Fixture framework + synthetic Tier-0 + shared replay helper** (also
delivered M2-core + M3 ? see below). `backend/eval/golden_fixtures.py` (`replay_fixture`, schema
dispatch, `Path(__file__)`-relative) + 3 **tool-generated** synthetic fixtures + manifest +
README + `.gitattributes` LF pin + `tests/test_golden_fixtures.py`. CONTROL-only (no numpy
execution on the replay path); cross-OS verified. | ? | No (synthetic) | ? |
[#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) |
153      | **M2** | **Freeze/refresh extractor.** Core (`freeze_synthetic` / `refresh_snapshot` /
```

```

`--check` CLI + idempotence test) shipped in M1 (#189); the `--license`/`--minors-
free`/`--owned` **refusal gate** shipped with P1 (this PR). | M1 | No | ? |
[#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) · this PR |
154 | **M3** | **Characterization gate** ? `tests/test_golden_fixtures.py`: parse-compared
(exact-int / approx-float 1e-6), **CONTROL-only**, positive no-person/audio assertion,
perturbation + idempotence. **Shipped in M1 (#189).** (Windowing/label fingerprint guard + row-
count cap ? fold into M4.) | M1 | No | ? | [#189](https://github.com/avidullu/badminton-
highlight-indexer/pull/189) |
155 | **MX** | **Cross-OS determinism guard (CI).** Lightweight `golden-crossos` CI job
(ubuntu/windows/macros) running `--check` + the fixtures test ? locks the cross-OS guarantee
(CONTROL replay is numpy-only, no GPU stack; a *separate fast job*, not full-suite×3). | M1 | No
| ? | [#191](https://github.com/avidullu/badminton-highlight-indexer/pull/191) |
156 | **M4** | **Quality-floor gate + first `regression_baseline.json`.
`tests/test_golden_quality_floor.py`; create `eval_baselines/regression_baseline.json` with
**windowing fingerprint**; metrics-vs-base paired-delta; corpus-level (not per-slug) LOVO for
runtime. | M1, M2 | No | ? | ? |
157 | **P1** | **De-attribution + minimization** ? `golden_fixtures.freeze_real_fixture()`
(extends the M1 module, not a separate script): **opaque random `fx_hex` surrogate slug**,
metadata strip, metric-invariant **time-origin reset**, strict GT loader, **positive no-
person/audio/MD5/source-path assertion**, manifest non-clobber; **--freeze-real` CLI with the
**refusal gate** (exit 2 unless owned + minors-free + license). Tooling-only, 12 synthetic
tests; **no real data committed**. | M1 | No (tooling) | ? | this PR |
158 | **P2** | **Owner-accepted (researched-risk), 2026-06-16.** Owner accepted the
researched IP/privacy analysis (§3) to proceed **under its conditions** (owned/Fork-A source ·
minors excluded · biometric/person streams excluded · de-attributed). Owner's informed risk
acceptance ? **NOT a retained-counsel opinion**; optional future counsel confirmation per §8. |
P1 | Owner-accepted | ? | ? |
159 | **P3** | **Tier-0 REAL subset ? SHIPPED.** Committed **6 curated, trimmed (?12k-frame
? few-min), de-attributed** real fixtures (`fx_r01`?`fx_r06`: single-court / multi-court×2 /
doubles×2 / continuously-active blob; both 29.97 & 59.94 fps) via `freeze_real_fixture --max-
frames` (opaque slug, time-origin reset, scenario-only labels, **NO** source name/path/video_id
? grep-verified + leak-scan clean). ~1.9 MB total; the characterization + cross-OS gates now run
on real data too. Owner cleared ALL owned golden videos (owned + minors-free) + re-confirmed
option (a) (2026-06-22). The surrogate?source map stays **off-git** (owner-held). | P1, P2 |
conditions §3 ? | ? | this PR |
160 | **O1** | ***(optional)* **Protobuf + deterministic serialization.** `rally.fixture.v1`
proto3, codegen + `protoc`/`git diff --exit-code` guard, canonical writer, **gitattributes`
binary. Only if cross-OS byte-stable snapshots are wanted. | M1 | No | ? | ? |
161 | **O2** | ***(optional)* **Tier-1 key-gated path.** `scripts/fetch_golden.py` (OneDrive
token, default) and/or SOPS+age; **_tier1_present()**; `RALLY_GOLDEN_DIR` wired; Tier-1 CI job
gated on secret, never on fork runners. | M1 | No | ? | ? |
162 | **DOC** | **Discoverability cross-links** ? added the `eval_baselines/fixtures/` row
to `CODE_MAP §0` + a `golden_fixtures.py` entry to `CODE_MAP §2.3`; cross-linked
`GOLDEN_DATA_SHARING.md` (committed subset ? private OneDrive corpus). README indexed in #186;
memory `current-state` kept live. (EVALUATION/QUALITY_ITERATIONS cross-links **deferred** ? that
doc surface is hot in the concurrent quality track; avoid collision.) | M1 | No | ? | this PR
|
163
164 **Must-fix-before-any-public-commit (red-team), folded into the rows above:** ? numeric-
tolerance not byte-equality + multi-OS matrix + CONTROL-only (M3/D5/MX ?); ? scrub names in
**all** files + history decision (P0); ? leak-scan a **required** check, hooks are convenience
only (P0); ? person-optional loader + **positive** no-person assertion, don't trust the silent
`0.0` default (P1 ?); ? random surrogate id not HMAC-over-roster (P1 ?); ? quantize/drop
fps/frame_width (P1 ? fps/fw retained as scale constants; quantize is a P3 option); ? document
Tier-0 omits the fusion path (DOC); ? counsel + Fork-A + minors/biometric exclusion as hard
preconditions (P2 owner-accepted; enforced in code by P1's refusal gate ?).
165
166 ### Progress log (quantitative ? coverage must hold/improve, CI incl. the 3-OS cross-OS
job green)
167 | PR | Deliverable(s) | Suite | Coverage | Notes |
168 |----|-----|-----|-----|-----|
169 | [#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) | M1 +
M2-core + M3 | 746 pass / 7 skip | 80.41% | synthetic Tier-0 + CONTROL replay +
characterization; cross-OS proven (CI Linux replayed windows-frozen fixtures) |
170 | [#191](https://github.com/avidullu/badminton-highlight-indexer/pull/191) | MX +

```

```

tracker/legal reconcile | full suite green | 80.41% | `golden-crossos` CI job green on
ubuntu+windows+macos ? cross-OS **enforced**; legal ? owner-accepted |
171 | _this PR | P1 + M2 refusal gate | 759 pass / 7 skip | **80.49%** (+0.08) |
`freeze_real_fixture` + refusal gate; **+13 tests**; tooling-only, no real data committed |
172 | _this PR (audit-hardening)_ | adversarial-audit fixes #17#17 | 784 pass / 7 skip |
**80.71%** (+0.22) | strict-slug `label` + dropped `source_note` (#1); 6dp-rounded reset within
`_FLOAT_TOL` (#2); temp-dir-then-promote + negative-label refusal (#3); exact-membership source-
path scan (#4); key-set-equality in `compare_expected` (#5); `--reason` mandatory (#6);
independent schema guard (#7); `max_win` recorded-not-applied note (#8); numpy-execution wording
(#9); +failure-semantic & by-kind tests (#10/#11); `test_golden_real_fixtures.py` added to
`golden-crossos` CI (#12); doc reconciled to shipped reality incl. per-frame-CSV honesty + risk
re-confirm (#13?#17). Synthetic `--check` byte-identical (unchanged). |
173 | _this PR (P0 safe scope)_ | personal-path scrub + leak-scan guard | 1365 pass / 7 skip
| 85.27% | scrubbed `C:\Users\?` paths from 3 code docstrings (`eval_partial_labels.py` /
`gemini_refine.py` / `README.md`); `tools/leak_scan.py` (32-hex MD5 + personal-path structural
guard + gitignored `leakscan_private` name-list) + a CI `leak-scan` job + 10 tests. Functional
video-name rename DEFERRED (owner safe-scope). |
174 | _this PR (P3 real fixtures)_ | 6 curated trimmed de-attributed real Tier-0 fixtures +
`--max-frames` trim | 1390 pass / 7 skip | 85.27% | `fx_r01`?`fx_r06` (~1.9 MB; single/multi-
court/doubles/blob, both fps) frozen via `freeze_real_fixture --max-frames=12000`; de-
attribution grep-verified + leak-scan clean; +trim test; owner re-confirmed option (a). F1
0.0?0.59 = real CONTROL-path (behaviour-locking, **NOT** field quality). |
175
176 ---
177
178 ## 8. Open owner / counsel questions
179

```

7. user (2026-07-02T03:00:08.084Z)

```

45
46 - ? **Annotation guideline** ? **DONE (2026-06-02):** authored as the Roora brief +
illustrated guide in the collector repo `docs/annotation/` ? rally CSV schema; ignore dead time
/ include every rally (incl. faults & lets) / shot = one contact; `ending_reason`
(forced/unforced-error taxonomy, shared across net-separated sports); singles & doubles for
badminton & pickleball.
47 - ? **Internal gold key** ? hand-annotate 5?10 licensed clips to near-perfect as the
vendor answer key. **Unblocked.** **Revisit:** anytime.
48 - ? **GS-3 ? `HumanVendorProvider`** ? concrete adapter for the chosen annotation vendor
(submit/poll/fetch + DPA/secure-link controls). **Vendor selected for the initial pilot: Roora**
(Bengaluru; iMerit / Taskmonk / TELUS remain the fallback shortlist). **Revisit:** once the
Roora pilot confirms quality, then build the adapter.
49 - ? **GS-4 oracle model** ? `AutomatedProvider` interface is stubbed (this work) with a
deterministic fake. Wire a `concrete` fast-tier auto-labeler later. **Oracle rule:** it MUST be
a different model lineage than production, or the regression test is blind to shared failures.
**Revisit:** when a production LLM step is chosen (pick the oracle's lineage to differ).
50 - ? **GS-5 ? regression triggers** ? wire `regress()` into CI (pre-merge gate), on-
demand CLI (the "surprise" case), and a scheduled run; include a deliberate re-baseline path.
**Revisit:** after GS-4.
51 - ? **Vendor pilot (Roora)** ? run the paid pilot with **Roora** as a **single-vendor
pilot first** (not a multi-vendor bake-off): assess quality-per-rupee against the $4 bar;
escalate to the iMerit/Taskmonk fallback only if it misses. **Revisit:** the pilot doubles as
calibration once the gold key exists.
52 - ? **Phase-1 video collection** ? gather the round-1 golden corpus per
`docs/GOLDEN_SET_PHASE1_VIDEOS.md`: **static self-recorded** footage (mirror deployment, not
broadcast), 3 sports incl. singles & doubles, across the diversity matrix; ~12 clips/sport with
**? held out by match**. **Unblocked.**
53 - ? **Train/eval separation guardrails** (was `DEFERRED_HARDENING_PLAN` #13; design
preserved here on archival 2026-06-14) ? train & eval videos live in **separate catalogs**
today, kept apart by convention only; nothing **enforces** that a video used to train the learned
segmenter (`backend/eval/training.py::build_training_set` ? `distill_local`) is absent from the
eval/regression set, so leakage would silently inflate the headline metrics. **Concrete design**
(from the archived DEFERRED plan $2): split unit = **match, never clip** (hold ~? of each sport
as eval, per `GOLDEN_SET_PHASE1_VIDEOS.md`); add a `split: "train"|"eval"` field to the
collector's `curate export` manifest (cross-repo contract, ADR-002); **[DONE ? pure-logic core,

```

```

PR #171]** `backend/eval/splits.py` now exists with `load_split(manifest) -> (train_ids,
eval_ids)` keyed by **match id** (not file MD5 ? re-encodes share a match but differ in
`compute_video_id`) + `assert_disjoint(train_ids, eval_ids)` that raises on overlap, with
`tests/test_splits.py` (overlapping ids ? `assert_disjoint` raises). **STILL DEFERRED:** wire
the guard into the training entrypoint (`distill_local`/`training`) and into `run_regression.py`
(baseline computed strictly on the eval split); add the collector's `curate export` manifest
`split` field (above); the integration test (a deliberately-leaked manifest fails the training
entrypoint). Full rationale/risks: [`archives/DEFERRED_HARDENING_PLAN-
COMPLETED-2026-06-14.md`](archives/DEFERRED_HARDENING_PLAN-COMPLETED-2026-06-14.md) §2.
**Revisit:** dedicated design session with the user (needs a collector-side manifest change).
54
55     ## Golden regression fixtures (video-free, non-attributable regression suite)
56     *(source: `docs/GOLDEN_REGRESSION_FIXTURES.md` §7 tracker / §9 definition-of-done)*
57
58     Core suite **shipped** (M1·M2·M3·MX·P1·P2·P3·DOC ? ? synthetic + 6 de-attributed real
`fx_r01`?`fx_r06` fixtures replay cross-OS, leak-guarded, CONTROL-only). The two items below are
the **gating remainder** of the doc's §9 definition-of-done (it flips to `DONE`/archive only
when **P0 + M1?M4 + P1 + DOC** are ?); 01/02 are explicitly optional. Tracked here so the track
isn't silently treated as "done" while DoD items remain.
59
60     - ? **M4 ? quality-floor gate + first `regression_baseline.json`** *(recommended-next;
unblocked).* Assert per-fixture **and** mean F1 ? a committed floor
(`tests/test_golden_quality_floor.py` + `eval_baselines/regression_baseline.json`). *First
move:* seed the per-slug floor from the F1 already frozen in each `fx_*/expected.json` `score`
(real fixtures sit at **0.0?0.59**, so floor = `current ? tolerance`, not a global bar).
**Gotchas:** record a **windowing fingerprint** in the baseline (else a SERVED retune silently
passes); keep it **CONTROL-only** (learned/numpy-BLAS is non-deterministic across CPUs); add it
to the `golden-crossos` CI matrix. **Revisit:** anytime ? it's the cheapest step to flip the
track to DONE.
61     - ? **P0 remainder ? functional video-name rename** *(deferred per owner safe-scope,
2026-06-16).* Rename the functional names (`eval_baselines/*.json` keys,
`rally_seg_eval.MULTICOURT_CLIPS`, quality-log names) ? opaque ids. *Why it's hard:* the keys
feed `training/gen0/harness.py`'s LOVO floor matched to the **gitignored
`output/human_lovo_manifest.json` + data files** ? rename them together or the floor breaks;
also collides with the hot quality track. *First move:* put the names in a gitignored
`.leakscan_private` and run `python -m tools.leak_scan --names .leakscan_private` to enumerate
every occurrence, then plan the coordinated rename. **Revisit:** before any open-sourcing /
public commit.
62     - ? **01 ? protobuf + deterministic serialization** *(optional).* `rally.fixture.v1`
proto3 + codegen + `protoc`/`git diff --exit-code` guard. **Revisit:** only if cross-OS **byte-
stable** snapshots are wanted (JSON parse-compare suffices today).
63     - ? **02 ? Tier-1 key-gated path** *(optional).* `scripts/fetch_golden.py` (OneDrive
token) and/or SOPS+age, `RALLY_GOLDEN_DIR` wired, Tier-1 CI gated on a secret (never on fork
runners). **Revisit:** only if sharing person/biometric streams (Tier-1) becomes necessary.
64
65     ## Recording best practices
66     *(source: `docs/VIDEO_RECORDING_GUIDELINES.md`)*
67
68     - ? **Validate recording hypotheses** ? empirically test which capture conditions
(resolution, fps, camera angle/height, lighting) actually move precision/recall, then promote
the winners into user-facing "how to record" guidance. Only the static-camera rule is currently
evidence-backed (ADR-004). **Revisit:** once a controlled corpus + the eval loop can measure
per-condition deltas.
69
70     ## Code & test organization
71     *(source: owner review comments on PR #35, 2026-06-07)*
72
73     - ? **`tests/` reorganization to reduce fragmentation** ? owner wants to isolate test
components on a grouping
74     metric (by subsystem/component) so `tests/` is less fragmented, mirroring the new
`backend/pipeline/audio/` +

```

8. user (2026-07-02T03:00:30.202Z)

origin <https://github.com/Khelsutra/badminton-highlight-indexer.git> (fetch)

```
origin https://github.com/Khelsutra/badminton-highlight-indexer.git (push)
{"isPrivate":true,"nameWithOwner":"Khelsutra/badminton-highlight-
indexer","visibility":"PRIVATE"}
```

9. user (2026-07-02T03:00:30.543Z)

tools/leak_scan.py

d51efe3 feat(training): Gen-0 TAL harness ? one-command train -> honest LOVO -> ship-gate ->
artifact (M0.4 scaffold) (#77)

10. user (2026-07-02T03:01:07.257Z)

Structured output provided successfully