

claude-windows / c1fcb4aa-537d-41fd-858e-53f93349bf5a

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\subagents\workflows\wf_7049e1ed-b93\agent-a80fcf97c631fce5a.jsonl`  
- SHA-256: `11fa38b7112d314485823ceb0436e9f43b042a2fa6a98e072b663f5c417004a3`  
- Source modified: `2026-07-02T04:57:13+00:00`  
- Imported at: `2026-07-05T16:48:24+00:00`  
- project: `wf_7049e1ed-b93`  
- session_id: `c1fcb4aa-537d-41fd-858e-53f93349bf5a`
```

Transcript

1. user (2026-07-02T04:45:39.414Z)

You are a senior engineer producing ONE isolated PR. Today is 2026-07-02.

REPO: C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector (a shared checkout ? follow the safety rules exactly).

AUDIT FINDINGS YOU ARE FIXING: R2-6 (golden-corpus registry rot) + critic #1 ? Grep C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/AUDIT_REPORT_khelsutra-guru_2026-07-02.md for these IDs and read the full evidence + verifier notes before coding.

TASK BRIEF:

In C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector: (1) add a 'curate verify' subcommand (or extend the existing curate CLI) that resolves every registered videos.local_path and reports dangling rows (non-zero exit with --strict); (2) re-point the 6 owned curated rows (adarsh_avi_singles, kushagra_singles, largetest_doubles, gbaaddy, testlarge_short, mahadevpura_singles) to where the files ACTUALLY exist now ? verify on disk first: check 'C:/Users/avidu/Projects/khelsutra-guru/Annotation Setup/' (Collect/Golden Labelled) and if the files are not on C: at all check the F: drive paths documented in C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/docs/DATA_IN_GCS.md \$5/\$6 ? point at the real existing location, and only for files you personally confirmed exist (Test-Path/ls); if some files exist nowhere reachable, leave those rows and flag them in the PR body; (3) the DB is a git-tracked binary sqlite ? commit the updated DB and note the binary-diff in the PR body; also purge or clearly mark the two 'watch?v=example' placeholder junk rows (critic finding) if trivially safe. Add tests for the verify logic (use a temp DB).

BRANCHING SAFETY (mandatory ? shared checkout, sibling sessions exist):

1. In the MAIN repo checkout run: git fetch origin
2. Create an ISOLATED WORKTREE: git worktree add "C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-<shortname>" -b <branch-name> origin/master (branch explicitly FROM origin/master ? never from the current local branch; use origin/main for repos whose default is main)
3. Do ALL work inside the worktree. Before committing: git branch --show-current + git log --oneline -1 must show YOUR branch on the origin-master base.
4. Stage EXPLICIT paths only (git add <file> <file>) ? NEVER git add -A / . / -u
5. Before opening the PR: git log --oneline origin/master..HEAD must list ONLY your commits.

COMMIT/PR CONVENTIONS:

- Commit message: conventional style matching the repo's history; end the message with a blank line then: Co-Authored-By: Claude Fable 5 <noreply@anthropic.com>
- PR body: what/why, a "## Verification" section stating exactly what you ran locally and what CI covers, and end with: ? Generated with [Claude Code](https://claude.com/claude-code)
- Do NOT merge the PR. Do NOT remove the worktree (a fixer may need it). Return the PR number, branch, worktree path.

QUALITY BAR:

- Read the repo's CLAUDE.md first; for the engine also docs/CODE_MAP.md \$0 (file placement) before adding any new file.
- Add/extend tests for every behavior change. Run the targeted tests you touched (python -m

pytest tests/<relevant> -q) + ruff check on changed files; if mypy config exists, run mypy on touched modules. Full suite is CI's job ? do not run it.

- Default-safe: new knobs default OFF / behavior-preserving unless the task explicitly says otherwise.
- The full audit evidence is in C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/AUDIT_REPORT_khelsutra-guru_2026-07-02.md ? Grep it for your finding IDs and READ the verifier notes; they contain corrections (exact file:line, scoping traps) you MUST honor.

If the brief's premise turns out false against current code (something already fixed, file moved), adapt honestly and say so in notes ? never force a pointless diff. If you genuinely cannot complete (env broken, auth missing), open NO PR and explain in notes with pr_number=0.

2. user (2026-07-02T04:45:50.668Z)

README.md
__pycache__
cleanup.py
config
data
docs
pyproject.toml
src
tests

3. user (2026-07-02T04:45:50.778Z)

486:### R2-6 [P1/M] ? sports-data-collector

4. user (2026-07-02T04:45:57.070Z)

```
440     ## Confirmed findings ? round 2 (8, previously unverified)
441
442     ### R2-1 [P2/S] ? rally-annotator
443     **Content script runs a 3 Hz full-DOM shadow-piercing scan on every page the user
visits, even with the panel hidden**
444     - verdict: CONFIRMED
445     - reasoning: Every element of the claim reproduces in the current working tree (rally-
annotator @ main 3d604bc, clean):
446
447     (1) Runs on every page: web/entrypoints/content.ts:13 `matches: ["<all_urls>"]` (top
frame only per line 16), and web/wxt.config.ts:13 `host_permissions: ["<all_urls>"]`.
content.ts:32-44 unconditionally calls `mountPanel(...)` then `panel.toggle()` ? the panel
starts hidden (display:none) but stays mounted.
448
449     (2) 3 Hz loop regardless of visibility: web/src/ui/panel.ts:376 `const clock =
window.setInterval(render, 333);` starts at mount and is cleared only in `destroy()`
(panel.ts:395), which is never called from production code ? only from
web/test/panel.smoke.test.ts. `toggle()` (panel.ts:391-393) merely flips display:none; it does
not pause the interval.
450
451     (3) Each tick triggers a full-DOM shadow-piercing scan on video-less pages: render()
calls `video.now()` (panel.ts:264) ? `active()` (web/src/video/directVideo.ts:55-58) ?
`findActive()` ? `collectVideos(document, ...)` (directVideo.ts:37-39), which does
`root.querySelectorAll("video")` plus `root.querySelectorAll("*")` with a per-element
`chrome.dom.openOrClosedShadowRoot` call, recursing into every shadow root to depth 8
(directVideo.ts:25-32, 13-23). Crucially, when no video exists, `this.current` is set to null
(directVideo.ts:40-43), so the cache NEVER fills and every 333 ms tick repeats the full scan.
Each tick also rewrites the hidden panel's DOM (status.innerHTML, list.replaceChildren,
panel.ts:265-281).
452
453     (4) No fix landed: git log for content.ts/panel.ts/directVideo.ts shows only ad6ac9c
(initial), 5ec1831 and 0a05ecf (i18n) ? no gating change. No
```

visibilitychange/IntersectionObserver/MutationObserver/activeTab anywhere in web/ (grep: only hit is the setInterval itself).

454

455 (5) Intent: ``<all_urls>`` is deliberate (DESIGN.md LD-7: panel hidden by default, icon is the explicit "annotate this page" gesture) ? but LD-7 gates only the UI, not the scan loop, so the perf cost contradicts the design's own "not intrusive" rationale.

456 - notes: Sharper failure scenario: user installs the extension and opens ordinary video-less pages (Gmail/Docs/GitHub ? tens of thousands of DOM nodes, many shadow roots). In every such tab, every 333 ms, forever, the content script walks the entire DOM (``querySelectorAll("*")``) and makes one `chrome.dom.openOrClosedShadowRoot` call per element ? continuous per-tab CPU burn and laptop battery drain, plus jank on heavy pages, with the panel never once opened. Cost scales with number of open tabs. On pages WITH a video the cache (`directVideo.ts:55-57`) makes ticks cheap, so the pathological case is precisely the "every other page" case the finder described.

457

458 Severity P2 not P1: no correctness/data-loss/security impact, and the extension is version 0.1.0/private (not store-published), so blast radius today is dev machines; would be P1 if shipped broadly.

459

460 Correctly-scoped fix (S): (a) don't start the 333 ms clock until the panel is first shown and pause/clear it when hidden ? gate it inside ``toggle()`` in `panel.ts`; (b) make render's clock read non-scanning (e.g. a ``peekNow()`` on `DirectVideoHandler` that uses only the cached element), keeping the full scan for explicit triggers (Refresh button, panel open). Optional follow-up (M, separate decision): replace the ``<all_urls>`` content script with `activeTab + scripting.executeScript` on icon click, which eliminates injection on every page entirely and drops the broad host permission ? aligns with LD-7's own rationale, worth raising but not required to fix the perf debt. Not OWNER_GATED ? pure code change, no spend or product pivot.

461

462 ### R2-2 [P2/S] ? rally-annotator

463 **Locked decision LD-i8 (web language persisted + seeded from browser language) is documented as shipped but not implemented ? panel always boots in English and forgets the choice on reload**

464 - verdict: CONFIRMED

465 - reasoning: Doc claims shipped: rally-annotator/docs/LOCALIZATION.md:55-57 (LD-i8: "Web: persisted in `chrome.storage.local` (seeded from the browser language; default en)") and :99-100 mark the web i18n phases done (PR #18 "selector", PR #19); web/README.md:22-23 says "persisted in-panel language selector". Code says otherwise: (1) web/src/i18n/index.ts:56 boots the module-level locale to `DEFAULT_LOCALE ("en")` ? a fresh content-script injection always starts English; (2) `negotiateLocale` (`index.ts:72-80`), which implements the `stored?navigator.languages?en` order, has ZERO runtime callers ? grep across web/ shows its only callers are unit tests (`web/test/i18n.test.ts:47-52`); (3) `panel.ts:18` declares `onLocaleChange` with the comment "optional: persist the choice (content script wires this)", but the actual `mountPanel` call in `web/entrypoints/content.ts:32-42` passes only `annotator/video/identity/download` ? `onLocaleChange` is never wired, so the change handler at `panel.ts:306-311` updates only the in-memory ``current``; (4) no storage read/write of any language key exists anywhere: `persist/store.ts` handles only "rally:"-prefixed rows, `background.ts` (all 27 lines) handles only downloads and icon clicks. Git: repo is at origin/main (3d604bc); no commit after 0a05ecf (#19) touches web locale persistence ? 3d604bc ("docs: sync READMEs, multi-agent audit fixes") actually re-affirmed the false "persisted" README claim. The E2E (`web/e2e/extension.spec.ts:80-101`) switches the locale within one page session and never reloads, so it can't catch this. Failure scenario: a Hindi rater (browser hi-IN) opens a video ? panel boots in English despite the "seeded from browser language" claim; they select ??????; on the next page load/navigation the panel is English again and the choice must be re-made on every video.

466 - notes: Claim verified exactly as stated; not OWNER_GATED (LD-i8 is already the locked decision ? this is pure implementation). Sharper framing: both halves of LD-i8-web are missing (no boot-time seeding AND no persistence), yet all plumbing already exists and is tested, which is why the fix is S: in `content.ts` `main()`, read e.g. ``lang`` from `browser.storage.local`, call `setLocale(negotiateLocale(stored, navigator.languages))` before `mountPanel`, and pass `onLocaleChange: (l) => void browser.storage.local.set({ lang: l })`; add a unit test for the wiring and extend the Playwright E2E with a reload-persists-locale assertion (the current per-locale E2E never reloads, which is the coverage gap that let the docs drift). Scope note: the VLC half of LD-i8 (config-file persistence, PRs #20/#21) was not part of this claim and was not verified either way. Severity P2: no data loss/crash, but a locked decision is documented and README-advertised as shipped while absent ? it defeats the localization value prop for the

target hi/kn/te raters (English boot every page load) and the doc-sync commit 3d604bc propagated the false claim.

467

468 ### R2-3 [P2/S] ? sports-data-collector

469 **Collector pins obsreport as a plain PyPI requirement while the name is unclaimed on PyPI ? dependency-confusion window and install broken today**

470 - verdict: CONFIRMED

471 - reasoning: (1) Bare PyPI-style pin exists in current tree:

C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector/pyproject.toml:38-40 ? `report = ["obsreport>=0.1.3"]`, with the comment at lines 32-37 itself admitting the intended form is a git+ssh direct reference. (2) Name is unclaimed on PyPI today (2026-07-02):

<https://pypi.org/pypi/obsreport/json> returns 404, and `pip index versions obsreport` returns "ERROR: No matching distribution found for obsreport" (personally observed). (3) The private package really is named `obsreport` (sports-obsreport/pyproject.toml:6, version 0.1.3, tag v0.1.3 exists) and lives at the private repo Khelsutra/sports-obsreport (gh api confirms private=true; avidullu/sports-obsreport redirects there). (4) Contrast confirmed: the engine already pins it safely ? badminton-highlight-indexer/pyproject.toml:68 `obsreport @ git+ssh://git@github.com/avidullu/sports-obsreport.git@v0.1.3`. (5) No fix landed: the bare pin was introduced in commit 038424d and no later commit touched it (git log -S on pyproject.toml); HEAD adcd206, working tree clean. Failure scenario A (today): `pip install .[report]` in any clean env fails resolution ? the extra is uninstallable unless the sibling checkout was pre-installed editable. Failure scenario B (window): anyone can register `obsreport` on PyPI and publish >=0.1.3; a later `pip install .[report]` then silently executes attacker code (sdist build hooks/import) instead of failing loudly.

472 - notes: Severity calibrated down from an implied P1: the exposure is real but narrow. `report` is an optional extra; the base install is unaffected; the collector treats obsreport as a soft dependency that no-ops when absent (sports-data-collector/src/reporting/__init__.py:15-22); and CI installs only `[dev]` (.github/workflows CI line 19), so no automated pipeline ever resolves the bare name ? the confusion window only fires when a human/agent runs `pip install .[report]` in a fresh env (in the documented sibling-editable workflow, pip sees obsreport 0.1.3 already satisfied and never queries PyPI). It is still a standing security window on a generic, squatable name, plus a broken-today extra, in a repo whose sibling project already solved the identical problem. Fix (S, one line): replace pyproject.toml:39 with `obsreport @ git+ssh://git@github.com/Khelsutra/sports-obsreport.git@v0.1.3` (prefer the canonical Khelsutra org path ? the avidullu path in the engine's pin only works via GitHub's transfer redirect and is itself mildly stale; consider updating badminton-highlight-indexer/pyproject.toml:68 too). setuptools accepts PEP 508 direct references without extra config (no hatchling-style allow-direct-references needed), and since sports-data-collector is not published to PyPI, the PyPI ban on direct references does not apply. Optional hardening (owner-gated, not required for the fix): claim the `obsreport` name on PyPI with a stub to close the squatting window permanently.

473

474 ### R2-4 [P2/S] ? badminton-highlight-indexer

475 **No timeout on any serving-path ffmpeg subprocess ? one hung ffmpeg permanently wedges the single-worker job queue**

476 - verdict: CONFIRMED

477 - reasoning: Single-worker executor: backend/api/job_worker.py:47-49 (max_workers=1) with the in-code "O2 RISK" admission at :32-39 that one hung job blocks all subsequent jobs and "the executor itself cannot enforce timeouts on the Thread". Serving-path ffmpeg/ffprobe calls with NO timeout=: (compile job, POST /api/compile main.py:1594 -> _execute_compile -> compiler.compile_highlights at main.py:1686, no wrapper) backend/pipeline/stitcher/compiler.py:130 (ffprobe), :321 (per-clip libx264 encode), :494 (concat); (process job, _execute_processing main.py:1132) validators backend/pipeline/validators/format.py:43, pts_continuity.py:35, resolution_fps.py:15; backend/utls/video.py:18,:75 (chunking); backend/pipeline/segmenters/gemini.py:166; backend/pipeline/audio/hit_detector.py:53. The mitigations listed in job_worker.py:34-37 do NOT cover ffmpeg: run_bounded wraps only the Gemini generate call (backend/providers/gemini.py:233, op_timeout_sec=180); detector timeout_sec covers only GPU runner subprocesses (detectors/base.py:116, native_wasb_runner.py:134). Recovery is restart-only: fail_stale_jobs is called once at startup (main.py:2030; database.py:1122-1125) and marks ALL queued/running jobs failed. No fix landed: git head e410136; commit 24178a7 (#374) documents the O2 risk but adds no timeouts. Secondary confirmed defect: /api/source proxy warm (main.py:657 daemon thread) runs annotation_proxy.py:81 ffmpeg with no timeout; a hang leaks the _proxy_inflight dedup entry forever (main.py:607-620), permanently disabling proxy generation for that video (does not wedge the job queue).

478 - notes: Claim is literally true, but "permanently wedges" should be "wedges until server restart, and the restart sweep (fail_stale_jobs) then fails every queued job". Sharper failure scenario: POST /api/compile for a video whose stored local filepath sits on a network-backed mount (Google Drive streaming / SMB) ? storage.acquire_local_input is identity for local paths (main.py:1683) ? the mount stalls mid-read (or the MP4 is truncated/corrupt and ffmpeg's demuxer loops); subprocess.run at compiler.py:321 blocks forever; the compile job stays 'running'; every subsequent /api/process and /api/compile returns 202 but stays 'queued' indefinitely with zero error surfaced (drains submitted at main.py:1635 queue behind the blocked task on the max_workers=1 executor); only a full restart recovers, at the cost of failing the whole queue. Severity P2 not P0/P1: the trigger class is rare (mitigations already bound the historically hang-prone ops ? Gemini generate 180s, GPU/WSL detector timeout_sec 1800s, Cloud Run poll max_wait_sec), the risk is explicitly documented in-code, and restart recovers; but impact when triggered is a silent total outage of the job pipeline. Fix (S, mechanical): add timeout= to the ~10 serving-path ffmpeg/ffprobe subprocess.run/check_output sites (config-driven, e.g. 30-60s for ffprobe probes, stitching-level 1800s default for encode/concat/chunking) and map subprocess.TimeoutExpired to the existing clean-job-failure paths (compiler already converts CalledProcessError to a failed-job result at main.py:1687-1695; annotation_proxy already catches broadly at annotation_proxy.py:82). The broader per-job wall-clock timeout job_worker.py:38 calls a "future slice" is a separate M item and not required to close this finding. Not OWNER_GATED ? timeout defaults are an engineering choice with no spend/product implications.

479

480 ### R2-5 [P1/S] ? badminton-highlight-indexer

481 **Engine CLAUDE.md 'Current state' section is ~3 weeks false ? claims platform/owned-model work 'has not started'**

482 - verdict: CONFIRMED

483 - reasoning: Personally reproduced against the working tree and git history. (1) C:\Users\avidu\Projects\khesutra-guru\badminton-highlight-indexer\CLAUDE.md:30-31 states "***No platform/owned-model implementation has started.**"; lines 34-37 say the "committed next PLATFORM slice = async job model" and the owned-model track "starts at a greenlit Phase-0 milestone... Do not start implementation without explicit owner approval"; lines 41-42 call `StorageBackend` + `ComputeBackend` registries "Future:". git blame shows lines 30-37 were last written 2026-06-10 (commit 62fb2fb) and never touched since, despite CLAUDE.md itself being edited 6+ times after (e735fd4 06-13 ... af1ccbe 06-30 ? all other sections). (2) Current-tree reality contradicts every one of those claims: backend/storage/ contains a StorageBackend ABC + local/gcs/gdrive/gdrive_api/s3/capabilities implementations; backend/compute/base.py + cloud_run.py exist (ComputeBackend); backend/api/job_worker.py + tests/test_async_jobs.py exist ? the "committed next slice" async job model already shipped (8bcf6ce #339 "make /api/compile async (202 + poll)", 2026-06-22); backend/flywheel.py + backend/worker + backend/eval/served_gate_a.py exist. Owned-model track: a555887 (2026-06-20) "M2 FINALIZED ? Gen-0 weights/0.1.0-l4 pushed"; 4134721 archives DECOUPLED_COMPUTE as "DELIVERED (M0-M2 + M3.5 on GCP)"; 334ad0f (#336) M11 data-flywheel; b72e6de (#319) training worker fix. 314 commits landed since 2026-06-13, through 2026-07-01 (e410136), including a recorded launch NO-GO decision 2026-06-30 (4a7bfbe #404) ? i.e., the project reached launch-consideration, the opposite of "not started". (3) No fix landed: git log -- CLAUDE.md shows the last commit touching it (af1ccbe, 2026-06-30) added only the bit-flip-PR convention + Python 3.10 line; the Current state section is untouched since 06-10. Bonus staleness in the same file: lines 18-19 claim rally_gate.py Gate A is "not wired into production" while backend/eval/served_gate_a.py and c8a493d "served Gate-A litmus" (#401) indicate a served Gate-A path now exists.

484 - notes: Sharper failure scenario: CLAUDE.md is auto-loaded project memory for every fresh agent session in this heavily agent-driven repo. A fresh session reads "No platform/owned-model implementation has started" + "Do not start implementation without explicit owner approval" and will (a) tell the owner the async job model is an un-started, gated future slice and propose designing/implementing what already exists (backend/api/job_worker.py), (b) treat StorageBackend/ComputeBackend as "Future" and re-design duplicates of backend/storage/ and backend/compute/, (c) misreport project status upward (e.g., in handoffs/audits) as pre-implementation when Gen-0 weights shipped and a launch NO-GO was formally recorded 06-30. This also directly violates the workspace's own "keep it honest ? reflects real, shipped state" convention. Correctly-scoped fix (S, one doc-only PR): rewrite CLAUDE.md "Current state" (lines 29-37) to reflect reality as of 2026-07-01 ? platform slice (async jobs, storage/compute backends, capabilities/assets endpoints) SHIPPED; owned-model M0-M2+M3.5 DELIVERED on GCP, Gen-0 weights pushed; launch NO-GO 2026-06-30 with pointer to #404/docs; fix line 41-42 ("Future" ? "implemented, see backend/storage/, backend/compute/"); fix lines 18-19 Gate-A "not wired" claim against backend/eval/served_gate_a.py; re-point "committed next slice" at whatever docs/NEXT_STEPS.md currently prioritizes. Not OWNER_GATED ? this is a factual truth-sync of docs

to shipped code, the same class as the repo's prior "audit truth-sync" commits (62fb2fb, 9736a3b); no spend or strategy decision required. Round-1 claim of "~3 weeks false" is accurate: section frozen 2026-06-10, contradicting implementation landed from ~2026-06-20 onward.

485

486 ### R2-6 [P1/M] ? sports-data-collector

487 **Golden-corpus registry rot: all 6 owned 'curated' videos point at a deleted absolute local_path and the 4.9GB source videos are single-copy with no documented backup**

488 - verdict: CONFIRMED

489 - reasoning: Registry rot half CONFIRMED by direct observation. sports-data-collector/data/sports_metadata.db (git-tracked; rows landed in commit 9e6f204 "#19 ingest the n=6 owned golden corpus", DB last touched 9f1e792 with no path fix since): videos table has 8 status='curated' rows; the 6 owned ones (adarsh_avi_singles, kushagra_singles, largetest_doubles, gbaaddy, testlarge_short, mahadevpura_singles; license_type=Proprietary, is_commercial=1) all carry local_path under 'C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled*.MP4' ? that directory does not exist. The folder moved to 'C:\Users\avidu\Projects\khelsutra-guru\Annotation Setup\' whose 'Golden Labelled\' now contains ONLY the six .rallies.csv label files; zero copies of the MP4s exist anywhere in the workspace (searched). The committed export artifact data/eval_export/manifest.json repeats the same 6 dead absolute paths. Code consequence at sports-data-collector/src/cli/curate.py:679-684: export of a curated video with missing local_path emits md5=None plus a warning, and src/core/manifests.py:23-26 documents md5 as "the indexer's join key" ? so any fresh export of the entire owned golden corpus loses the deterministic MD5 join to badminton-highlight-indexer. Data-loss half CONFIRMED as single-copy but the "no documented backup" wording is wrong: (a) verified live via read-only 'gcloud storage ls' (2026-07-02): gs://khelsutra-rally-corpus/videos/ has 10 objects and NONE of the six originals, while labels/ DOES hold all six (2026-06-08_adarsh_avi_singles.rallies.csv etc.) and trajectories are in the bucket ? annotation work-product is safe, source footage is not; (b) Google Drive golden folder (G:\My Drive\Sharing\Annotation Videos\Golden Label Videos Request - June 15) holds only the 7 proxy clips, not the original six; (c) the sole remaining copy is documented at F:\[Khelsutra] GoPro Hero Black 12\KhelSutra\Training\Golden Labelled\ per badminton-highlight-indexer/scratch/copy_videos_to_gcs.ps1:31-36 and docs/DATA_IN_GCS.md:153 (\$6 Phase-2) ? and the F: drive is NOT currently mounted (only C: and G: exist), so even that copy's integrity is unverifiable today, and no MD5s of the six were ever persisted anywhere (eval_export/manifest.json predates the md5 field from PR #51; no proxy-registry rows).

490 - notes: Corrections to the round-1 claim: (1) Size is ~37.9 GB, not 4.9 GB ? per-clip sizes annotated in scratch/copy_videos_to_gcs.ps1:31-36 (10.9+7.7+11+6.7+1.4+0.19 GB); the ~4.9-5.2 GB figure matches the three unrelated Roora pilot videos under 'Annotation Setup/Archive/.../Annotation Videos/'. (2) "No documented backup" is false ? badminton-highlight-indexer/docs/DATA_IN_GCS.md \$5/\$6/\$8 explicitly documents the F: archive location and a Phase-2 migration plan (tracker at doc line 246: 14/15 videos still unmigrated); but "single-copy" stands: exactly one copy of the six source videos exists, on an intermittently-mounted external HDD, unverifiable right now, with no persisted hashes. (3) Downgrade from potential P0 to P1 because labels + trajectories for all six ARE in gs://khelsutra-rally-corpus (verified), so GPU-free eval/litmus and synth-train survive a drive failure; what would be permanently lost is the proprietary source footage (no re-extraction via --trajectory-source detector, no re-annotation, no re-rendering). Sharper failure scenarios: (a) F: HDD failure = permanent loss of ~38 GB of irreplaceable owned golden footage; (b) today, 'curate export' on the owned corpus emits md5=None for all six (curate.py:679-684), silently breaking the indexer MD5 join contract added in PR #51. Correctly-scoped fix (M, cross-repo): mount F:, verify + hash the six MP4s (persist MD5s into the DB/manifest immediately ? S and decision-free), generate 1080p proxies, upload via the existing scratch/copy_videos_to_gcs.ps1 or the rclone cloud path (Phase-2, already planned), then reimport/annotate local_path (or add a gs:// video_uri column) in sports_metadata.db and refresh data/eval_export/manifest.json. Partially OWNER_GATED sub-decision only: DATA_IN_GCS.md \$7.4 leaves proxies-only vs full-4K migration scope OPEN and the upload is real (though tiny, cents/month) GCS spend; the hash-and-record step and the DB path fix need no owner decision.

491

492 ### R2-7 [P2/S] ? non-git

493 **Roora vendor pilot record (paid deliverables + corrected golden labels + 3 videos) is a single unversioned local copy at risk of loss**

494 - verdict: CONFIRMED

495 - reasoning: CONFIRMED for the paid deliverables + corrected golden labels; the "3 videos" component of the claim is REFUTED (Drive copies exist). Evidence: (1) Single-copy artifacts verified on disk: raw paid CVAT deliverables at 'C:\Users\avidu\Projects\khelsutra-guru\Annotation Setup/Archive/Roora Pilot Review - 6 June 2026/' ? badminton/{badminton,

badminton (1), badminton (2)}/Annotations/*.xml (~2.3 MB per-frame boxes), pickleball/*.zip (3 zips, 260 KB), tabletennis/*.zip (3 zips, 213 KB) ? plus the human-corrected golden labels 'Annotation Setup/golden_labels_badminton.csv' (2.5 KB, mtime 2026-06-06) and the rescued multisport goldens 'Annotation Setup/Trajectories/roora/{pickleball,tabletennis}_rallies.csv'. (2) Not in git: sports-data-collector tracks only the AS-DELIVERED parse (data/roora_badminton_canonical.csv, verified via git ls-files + diff ? it still has rally 2 = 0:02.503?0:32.533, rally 7 end 1:51.612, no rallies 7.5/16.5), while golden_labels_badminton.csv holds the corrections (rally 2 start?0:29, rally 7 end?1:38, missed rallies 7.5@1:46 and 16.5@3:34 added, winner?out fixes) ? those corrections exist NOWHERE else (find over C:/Users/avidu/Projects + Downloads returned only this one file). (3) Not in GCS: recursive 'gcloud storage ls -r gs://khelsutra-rally-corpus/**' contains zero objects matching roora|pilot|pickle|tabletennis|sample_long|golden_labels; labels/ holds exactly the 15 canonical corpus clips (pilot video not among them); archive/ holds only a droplet archive. (4) Not on Drive: 'G:/My Drive/Sharing' has guidelines/briefs and 'Golden Annotations' (MahadevpuraSingles only) ? no CVAT exports. (5) Outside every backup convention: badminton-highlight-indexer/docs/DATA_IN_GCS.md §5 gap table (lines 117?131) covers only the 15-clip corpus and never mentions the pilot artifacts (and its line 121/142 path 'C:\?Annotation Setup\Collect\Trajectories\' is stale ? the folder moved, no Collect dir exists); the F: backup (memory checkpoint 2026-06-22) covered only the 7 new-clip labels; docs/archives/roora-vendor/README.md archives only guidelines+ADR, not pilot results; F: is currently unmounted. REFUTED part: the 3 pilot videos have byte-size-identical cloud copies at 'G:/My Drive/Sharing/Annotation Videos/Pilot Videos/' (Sample Long Video - Baadminton.MP4 = 2,882,861,900 B; Video 2 - Pickleball.mp4 = 1,435,276,306 B; Video 3 - Tabletennis.mp4 = 871,944,715 B ? all match local), and the local 'Annotation Videos/desktop.ini' references GoogleDriveFS.exe proving Drive File Stream provenance. Owner's own memory audit had flagged the precursor ("Roora multisport CSVs live ONLY in %TEMP%? one cleanup from loss"); the rescue moved them to Annotation Setup but they remain a single unversioned C:-disk copy.

496 - notes: Correction to the finding: the at-risk payload is NOT the 4.9 GB folder ? the videos are safe on the owner's Google Drive. What is genuinely one C:-disk failure (or one folder-cleanup) from loss is ~3 MB of paid, partly irreplaceable work product: the raw vendor CVAT exports for all 3 sports (box geometry the tracked canonical CSVs discard), the human frame-by-frame correction record (golden_labels_badminton.csv + rally_review.csv ? hours of review labor and the calibration evidence behind the vendor feedback), the pickleball/tabletennis golden CSVs backing the recorded TT-0.909/PB-0.308 multisport result, and the pilot review MD/PDF record. Vendor re-export from CVAT job 4075945 is not under owner control. Failure scenario: C: SSD failure or a cleanup of the non-git 'Annotation Setup' folder silently destroys the only record of the paid pilot and its corrected labels; nothing in DATA_IN_GCS.md's migration plan would ever pick it up because the pilot artifacts are absent from its §5 gap table. Fix (S, not owner-gated for the metadata): one 'gcloud storage cp -r' of 'Annotation Setup/Archive/Roora Pilot Review - 6 June 2026/' (excluding Annotation Videos/, already on Drive) + 'Annotation Setup/{golden_labels_badminton.csv,rally_review.csv,Trajectories/roora/}' to gs://khelsutra-rally-corpus/archive/roora-pilot-2026-06/ (archive/ prefix precedent already exists; ~3?10 MB, negligible cost), then add a Roora-pilot row to DATA_IN_GCS.md §5 and fix its stale 'Annotation Setup\Collect\Trajectories' path (lines 121/142). Including the 3 pilot videos in the bucket is the only owner-gated piece (DATA_IN_GCS §7 open decision 4, ~5 GB ? \$0.10/mo) and is optional given the Drive copies.

497

498 ### R2-8 [P1/S] ? badminton-highlight-indexer

499 **Upload retention sweep tool exists but nothing schedules it ? unbounded upload growth + unenforced 7-day privacy retention on any hosted deployment (#301)**

500 - verdict: CONFIRMED

501 - reasoning: Tool exists: badminton-highlight-indexer/tools/cleanup_caches.py (T11 retention sweep, --uploads-retention-days default 7.0 at ~line 407; classify() purges uploaded videos + .partial-* in the uploads zone; landed cb7394b/#299). Nothing schedules it: workspace-wide grep for cleanup_caches hits only the tool/tests/docs and an error string (backend/pipeline/detectors/base.py:162); zero references in .github/workflows/nightly.yml, deploy/cloudrun/, deploy/digitalocean/, scripts/, and zero cron/timer/sweep references in khelsutra/deploy/droplet/ (khelsutra-engine.service has no companion timer). Backend startup only sweeps Gemini remote orphans (backend/main.py:1808 _maybe_sweep_gemini_orphans), not local uploads. Issue #301 is OPEN (viewed via gh today); no scheduling commit in git log between cb7394b and master e410136. Deployment is LIVE: khelsutra/deploy/DEPLOY_ALPHA.md line 3 "Status: LIVE (deploy live-verified on the droplet 2026-06-22)"; uploads land via backend/api/uploads.py into security.upload_dir (/srv/khelsutra/incoming/uploads per runbook §2.2).

502 - notes: Sharper failure scenario: the alpha droplet is live and accepting uploads

behind Cloudflare Access; without any scheduled `python -m tools.cleanup\_caches --apply`, a tester's footage uploaded day 0 is still on disk day 30+ (unenforced 7-day privacy retention, an explicit owner policy per the tool docstring, 2026-06-21), and output/uploads grows unboundedly until require_free_bytes (backend/api/uploads.py:93) starts rejecting new uploads (service degradation on the small CPU droplet). Mitigations that keep this P1 not P0: runbook notes testers had not yet exercised the upload CUJ (OTP login + Gemini key rotation pending as of 2026-06-22), and the owner can run the sweep manually. Fix scope (S): do NOT re-implement the sweep ? ship a systemd timer + oneshot unit (or cron line) alongside khelsutra/deploy/droplet/khelsutra-engine.service that runs `python -m tools.cleanup\_caches --apply --uploads-retention-days 7 --uploads-dir <security.upload\_dir>` daily with --json output captured to journald; default-OFF/opt-in flag so local single-user installs are unaffected; document enable+verify steps in DEPLOY_ALPHA.md; one smoke test asserting the wrapper passes the right flags (tool logic already unit-tested in tests/test_cleanup_caches.py). Issue #301's acceptance criteria already spec this exactly. Partial owner gate: shipping the timer unit + docs is not gated, but *enabling* it on the live droplet is an ops action requiring the owner's SSH access, and #301 flags one open design call (whether apply should also skip uploads referenced by non-terminal jobs rows) ? recommend defaulting to the time-window-only behavior and letting the owner ratify.

503

504

505 ## Owner-gated (flagged, no action)

506

507 - [P2] badminton-highlight-indexer: **P0-rename remainder: person/venue-attributable video names still committed in code and baselines** ? Defect verified in current tree: backend/eval/rally_seg_eval.py:52-54 commits MULTICOURT_CLIPS = frozenset({"mahadevpura_2", "GX010128", "Badminton_BXH_2", "Boxhill_Doubles"}) (venue-attributable), and eval_baselines/heuristic_lovo_n6.json per_video keys are adarsh_avi_singles, kushagra_singles, largetest_doubles, gbaaddy, testlarge_short, mahadevpura_singles (person first names + venues). git log on rally_seg_eval.py (latest 0c9c693 #395) and on the baseline (#77) shows no rename ever landed. docs/GOLDEN_REGRESSION_FIXTURES.md:151 and docs/BACKLOG.md:61 track this exact remainder as DEFERRED per the owner's safe-scope decision (2026-06-16) with trigger "before any open-sourcing / public commit" ? a product/strategy owner gate ? and the coordinated fix requires renaming the owner-held gitignored output/human_lovo_manifest.json in lockstep or the gen0 LOVO training floor breaks. Repo is currently PRIVATE (gh repo view: Khelsutra/badminton-highlight-indexer, isPrivate=true), so no live public exposure today. Real, unfixed debt, but remediation is explicitly owner-gated on the open-sourcing decision and owner-side data coordination.

508

509 - [P2] khelsutra (avidullu/khelsutra): **Issue #114 stale-open: CF Access guest-login debugging against a deployment that no longer exists** ? Debt is real but the fix is an owner tracker call on the owner-gated alpha/serving track. Verified: (1) gh issue view 114 ? OPEN, createdAt == updatedAt == 2026-06-23T10:25:52Z, 0 comments; body is CF-dashboard-only triage of the Access allowlist for the droplet alpha and explicitly says "Deferred by owner 2026-06-23". (2) Cited anchor exists: deploy/DEPLOY_ALPHA.md:106 is the Allow-policy/Emails-allowlist step (line 105 is the app-domain step ? off by one). (3) Deployment gone: memory session-handoff.md:19 ? droplet claude-do-rally-test/168.144.126.176 deleted 2026-06-25, "the prior 'LIVE alpha on the droplet' ... is GONE ? alpha is paused"; api.khelsutra.guru CNAME/CF-Access flagged for cleanup. (4) Not already fixed: git log -- deploy/DEPLOY_ALPHA.md shows nothing relevant after 2026-06-23 (latest 6e4dbaf PR #113 is the unrelated GPU-VM doc); no in-repo reference to #114. (5) Nothing in-repo can fix or verify it ? the Access policy lives in the Cloudflare dashboard and the origin behind app./api. no longer exists. Fix requires owner disposition (close vs retitle), squarely under the alpha/serving go-live gate.

510

511

512 ## Critic additions (round 1 completeness pass)

513

514 - [P1/M] critical-fix: **Golden-corpus registry rot: all 6 owned 'curated' videos point at a deleted absolute path, and the 4.9GB source videos are single-copy with no documented backup** (`sports-data-collector/data/sports\_metadata.db:0`)

515 - evidence: Query of the git-tracked registry: videos rows adarsh_avi_singles, kushagra_singles, largetest_doubles, gbaaddy, testlarge_short, mahadevpura_singles (all is_commercial=1, status=curated) have local_path='C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\...' and video_url=NULL. `ls /c/Users/avidu/Projects/Annotation Setup` -> 'No such file or directory'; the data now lives at 'khelsutra-guru/Annotation Setup/Golden Labelled' (folder moved, registry never updated). src/cli/curate.py:224 registers

owned videos with `video\_url=None, local\_path=video\_path` ? raw operator-supplied absolute path, no normalization or later integrity check. ShuttleSet rows are also CWD-relative with mixed separators (`.data/videos/shuttleaset\_1.mp4`). 'Annotation Setup/Archive' is 4.9G of source match videos; grep for gs://|bucket|backup across sports-data-collector docs/README returns nothing, while docs/MANUAL_INGESTION_POSITIONING.md:13 declares the collector the 'single canonical system-of-record' for golden rally labels.

516 - fix: Add a `curate verify` (or extend cleanup.py) that resolves every registered local_path and flags dangling rows; re-point the 6 golden rows to the current 'khelsutra-guru/Annotation Setup/Golden Labelled' location; store paths relative to a configurable corpus root instead of absolute machine paths; and back the 4.9GB Archive up to the existing GCS bucket, recording the gs:// URI in video_url so the registry has a durable pointer.

517

518 - [P2/S] ci-infra: **sports-obsreport has no CI at all ? combined with the engine's skipped reporting tests, zero reporting-stack code is gated anywhere** (`sports-obsreport/pyproject.toml:7`)

519 - evidence: `ls sports-obsreport/.github/workflows` -> no such directory (every other repo in the umbrella has workflows). The repo has 10 test suites (tests/test_customer_message.py ... test_spec.py) and a dev extra with pytest+pytest-cov, but nothing runs them on push/PR to github.com/Khelsutra/sports-obsreport (HEAD = v0.1.3, main). The engine consumes it as a tagged dependency (badminton-highlight-indexer/pyproject.toml:68) and the engine's own 17 obsreport-gated tests never run in CI (already-confirmed finding) ? so no test anywhere gates the reporting stack.

520 - fix: Add a minimal .github/workflows/ci.yml to sports-obsreport: checkout, pip install -e .[dev], pytest --cov on push/PR. This closes the library half of the reporting-stack gap independently of the engine's obsreport-lane work.

521

522 - [P3/S] hygiene: **Engine .git bloated to 281MB by ~229MB of dangling raw match-video blobs; 6.5MB yolov8n.pt is permanently in pushed history and .gitignore has no guard against re-staging raw videos** (`badminton-highlight-indexer/.gitignore:49`)

523 - evidence: `git count-objects -vH`: 4716 loose objects, 266.33 MiB loose vs only 2.61 MiB in packs; du .git = 281M. Four loose blobs are raw MP4s (ftypisom headers): 22.4MB, 31.7MB, 84.5MB, 90.2MB ? all unreachable even via `git rev-list --all --reflog` (dangling from an aborted `git add`), never gc-pruned. Separately, yolov8n.pt (6,549,796 bytes) is reachable history on origin/master (added 5257fd5, removed d19557e 'replace AGPL ultralytics YOLO'), so every clone carries the AGPL-era weights forever. .gitignore covers *.pt (line 5) but only two narrow mp4 patterns (lines 49-50: backend/static/highlights_*.mp4, *_compiled.mp4) ? nothing stops the raw-match-video staging accident from recurring.

524 - fix: Run `git gc --prune=now` in the local checkout to reclaim ~260MB, and add a broad guard (e.g. `\*.mp4` with explicit `!` exceptions for any tracked fixtures, or ignore the video scratch dirs) to .gitignore. Optionally note yolov8n.pt in docs as accepted history weight; rewriting pushed history is not worth it at 6.5MB.

525

526 - [P3/S] tech-debt: **obsreport dependency URL still points at the pre-transfer avidullu org ? installs survive only via GitHub's transfer redirect** (`badminton-highlight-indexer/pyproject.toml:68`)

527 - evidence: pyproject.toml:68: "obsreport @ git+ssh://git@github.com/avidullu/sports-obsreport.git@v0.1.3" (echoed in requirements.txt:35). The repo actually lives at Khelsutra/sports-obsreport ? `gh repo view avidullu/sports-obsreport` resolves to nameWithOwner Khelsutra/sports-obsreport, i.e. a transfer redirect. The redirect breaks the moment any repo is re-created under the old name, and like the already-flagged trackers pin this is a mutable tag, not a commit SHA. The local clone's own remote already uses the Khelsutra URL, so the pyproject reference is the only stale spot.

528 - fix: Update the dependency URL in pyproject.toml (and the requirements.txt comment) to github.com/Khelsutra/sports-obsreport.git and pin the v0.1.3 commit SHA instead of the tag, matching whatever convention the trackers-pin fix adopts.

529

530 - [P3/S] hygiene: **1.7GB of stale yt-dlp partial-download debris in data/videos and placeholder junk rows in the committed corpus registry** (`sports-data-collector/src/crawlers/badminton\_crawler.py:231`)

531 - evidence: data/videos contains _diag.mp4.part (354MB), _t.mp4.part (516MB), _v.mp4.part (849MB) plus Frag*.part and .ytdl markers, all dated 2026-05-26 ? 1.7G total of dead partial downloads (du -ch *.part *.ytdl). The crawler's failure cleanup (badminton_crawler.py:231-246) removes only `local\_path`, never yt-dlp's .part/.part-Frag/*.ytdl side files. The tracked sports_metadata.db also carries two placeholder rows with video_url='https://youtube.com/watch?v=example' (staging_match_non_comm, shuttleaset_12) ? test

fixtures committed into the canonical registry.

532 - fix: Delete the stale *.part/*.ytdl files; extend the crawler's failure cleanup (or
cleanup.py) to glob and remove yt-dlp fragment side files for the failed target; purge or
clearly mark the two 'watch?v=example' placeholder rows out of the committed registry.

533

534

535 ## Unverified backlog (dropped at the round-1 cap; lower-ranked or duplicate)

536

537 - [P2] sports-data-collector: Stale doc pointers after archive moves (cross-repo
DECISIONS.md path + 6 in-code CODE_HEALTH_PLAN references)

538 - [P2] sports-data-collector: curate.py is a 1531-line monolith (CLI + interactive
review + import/export/proxy/licensing)

539 - [P2] sports-data-collector: Golden-label store is a binary SQLite file tracked in git
? unmergeable across the owner's concurrent sessions

540 - [P1] rally-annotator (Khelsutra/rally-annotator; local checkout
C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): Content script runs a 3 Hz full-DOM
shadow-piercing scan on every page the user visits, even with the panel hidden

541 - [P1] rally-annotator (Khelsutra/rally-annotator; local checkout
C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): Locked decision LD-i8 (web language
persisted + seeded from browser language) is documented as shipped but not implemented ? panel
always boots in English and forgets the choice on reload

542 - [P2] rally-annotator (Khelsutra/rally-annotator; local checkout
C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): Issue #26 open and unaddressed: CSV
carries zero provenance ? labels are unattributable to their source video except by filename

543 - [P2] rally-annotator (Khelsutra/rally-annotator; local checkout
C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): E2E never reloads a page: the two
headline persistence behaviors (resume numbering from storage, locale persistence) have no end-
to-end proof, and entrypoints/ are outside unit coverage

544 - [P2] rally-annotator (Khelsutra/rally-annotator; local checkout
C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): Issue #22 still fully pending: all six
translation catalogs remain machine-draft ? the declared quality gate before any public launch

545 - [P2] rally-annotator (Khelsutra/rally-annotator; local checkout
C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): 'Annotation Setup' is the un-versioned
annotation-data workspace (not a code duplicate) and its Roora pilot-review artifacts appear to
exist nowhere else

546 - [P1] sports-obsreport: Collector pins obsreport as plain PyPI requirement while the
name is unclaimed on PyPI (dependency-confusion window; install broken today)

547 - [P2] sports-obsreport: Engine (and both repos' comments) reference obsreport via the
stale pre-transfer org URL avidullu/ instead of Khelsutra/

548 - [P2] sports-obsreport: sports-obsreport has zero CI ? tags/pushes are cut with no
automated test or lint gate

549 - [P2] sports-obsreport: Engine-obsreport integration is never exercised in any CI ? the
~12 reporting tests only ever run on dev machines

550 - [P1] non-git dirs: "Annotation Setup" + "khelsutra-screencast" (workspace root
C:/Users/avidu/Projects/khelsutra-guru): Roora vendor pilot record (paid deliverables +
corrected golden labels + 3 videos) is a single unversioned local copy

551 - [P2] non-git dirs: "Annotation Setup" + "khelsutra-screencast" (workspace root
C:/Users/avidu/Projects/khelsutra-guru): khelsutra-screencast: 8 ad-hoc QA scripts and doc-cited
DoD evidence are unversioned; the 5 cited checkpoint screencasts are absent from the canonical
dir

552 - [P2] non-git dirs: "Annotation Setup" + "khelsutra-screencast" (workspace root
C:/Users/avidu/Projects/khelsutra-guru): DATA_IN_GCS.md §5 gap table contradicts its own §7:
trajectories/labels marked 'not in bucket' though migrated 2026-06-22

553 - [P2] non-git dirs: "Annotation Setup" + "khelsutra-screencast" (workspace root
C:/Users/avidu/Projects/khelsutra-guru): Stale-open #175: OneDrive golden-data sharing plan
never executed and largely superseded by the GCS bucket

554 - [P1] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra,
sports-data-collector, rally-annotator, sports-obsreport: Vendored SPA bundle in engine is 10
web/ commits stale ? engine-served UI lacks the fps display its own API just shipped

555 - [P1] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra,
sports-data-collector, rally-annotator, sports-obsreport: Engine PR #390 (D1 main.py split) is
now CONFLICTING against master and going staler

556 - [P1] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra,
sports-data-collector, rally-annotator, sports-obsreport: Engine CLAUDE.md 'Current state' is ~3
weeks false ? claims platform/owned-model work 'has not started'

557 - [P1] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: D0 droplet retired 2026-06-25 but runbooks in BOTH repos still say the deploy is LIVE on it

558 - [P2] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: sports-obsreport has zero CI ? no .github directory at all despite shipping versioned releases

559 - [P2] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: Lint/type gates exist only in the engine ? sports-data-collector CI is pytest-only, plus GH-Actions version drift across repos

560 - [P2] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: sports-data-collector working tree is dirty: uncommitted .gitignore WIP + a full stale nested clone of itself

561 - [P0] badminton-highlight-indexer (Khelsutra): #340/#330: /api/process still re-bills Gemini on re-submit of an already-processed video ? no idempotency short-circuit (live P0 cost bug)

562 - [P1] badminton-highlight-indexer (Khelsutra): PR #390 (D1 main.py split, step 1) is CONFLICTING and stale ? master diverged 9 commits on main.py since it was cut

563 - [P1] badminton-highlight-indexer (Khelsutra): #332: NSFW pre-flight gate still absent ? non-owner alpha users can route arbitrary video to Gemini (ToS/abuse + paid-garbage risk)

564 - [P2] badminton-highlight-indexer (Khelsutra): #327: path-jail still blocks gs://-style ingest in hosted mode ? validate_input_path runs before scheme resolution

565 - [P2] badminton-highlight-indexer (Khelsutra): #388 premise is half-stale: the three job-worker functions already live in backend/api/job_worker.py (since PR #234, 06-18) ? remaining scope is shim removal + test patch-target migration

566 - [P2] badminton-highlight-indexer (Khelsutra): #384: coverage floor 70 and offline guard lanes still duplicated verbatim between ci.yml and nightly.yml

567 - [P2] badminton-highlight-indexer (Khelsutra): #301: upload retention sweep tool exists but nothing schedules it ? unbounded upload growth on any hosted deployment

568 - [P2] badminton-highlight-indexer (Khelsutra): #349 (+#331 overlap): WASB streaming decode still forced single-threaded ? L4 GPU ~80% idle; #331 additionally needs a refresh (droplet baseline retired, its lever 3 already shipped)

569 - [P2] badminton-highlight-indexer (Khelsutra): #328: per-video runlog/trace endpoint still missing ? obsreport exists server-side but is not user-accessible

570 - [P1] Khelsutra/badminton-highlight-indexer: No timeout on any serving-path ffmpeg subprocess ? one hung ffmpeg permanently wedges the single-worker job queue

571 - [P1] Khelsutra/badminton-highlight-indexer: MD5 segmentation reuse still unimplemented ? every re-submit of a byte-identical video re-bills Gemini and re-pays the full CPU pipeline

572 - [P1] Khelsutra/badminton-highlight-indexer: WASB streaming decode still single-threaded (num_workers forced to 0), no NVDEC anywhere ? L4 GPU ~80% idle during detection

573 - [P1] Khelsutra/badminton-highlight-indexer: Upload retention sweep still manual ? 7-day privacy retention window unenforced without an operator (#301 open)

574 - [P2] Khelsutra/badminton-highlight-indexer: Remote-input scratch copies (rally-input-*/rally-worker-*) are never reclaimed ? by-design temp leak outside all cleanup tooling

575

576

577 ## Unverified P3 pool

578

579 - badminton-highlight-indexer: regression.py silently passes when no baseline exists (audit B8, still open) (critical-fix, `backend/eval/regression.py`)

5. assistant (2026-07-02T04:46:02.800Z)

Now let me read the repo docs and inspect the current state.

6. user (2026-07-02T04:46:03.737Z)

```
1      # Sports Data Collector & Curator
2
3      A clean, extensible Python codebase and CLI toolchain designed to discover, ingest,
validate, and curate sports video datasets. The repository stores annotations and video metadata
in an SQLite database, supports sport-specific tables and schemas, manages video storage outside
of Git, and enforces strict commercial license compliance.
4
5      > **Pipeline role:** this repo is the *data producer* feeding the **badminton-highlight-
indexer**, which trains/evaluates an offline rally segmenter on this curated, commercially-
```

licensed data. See docs/PIPELINE_ROLE.md for why `export`/`fetch` exist and how the two repos fit together.

```
6
7    ---
8
9    ## Features
10
11    - SQLite Backend: Centralized `sports_metadata.db` with a generic `videos` table and
12      sport-specific annotation tables (`badminton_rallies`, `tennis_points`, `cricket_deliveries`,
13      `pickleball_rallies`, `tabletennis_rallies`).
14    - Commercial License Flagging: Each ingested source license is checked against a
15      configured whitelist (`MIT`, `Apache-2.0`, `BSD`, `CC0`, `CC-BY`, `Public-Domain`,
16      `Proprietary`) and the resulting `is_commercial` flag is stored alongside the video. Non-
17      commercial sources (like `CC-BY-NC-SA`) are flagged as non-commercial but still ingested ?
18      nothing is discarded automatically. The whitelist is read from `config/config.yaml` consistently
19      across the parser, validator, and CLI.
20    - Overlap & Duration Validation: Ensures chronological event bounds (e.g. rally
21      start/end times) do not overlap and that durations are positive.
22    - Stroke-level CSV Parser: Ingests ShuttleSet/CoachAI-style stroke CSV logs and
23      aggregates them into rally durations using configurable start/end time buffers.
24    - Local Ingestion Mode: Manually registers local videos and associated CSV/JSON
25      annotation files directly into the database as curated entries.
26    - Interactive Curation CLI & Visual Play Harness: Let's you query DB status, list
27      staging entries, review sample timestamps, and approve/reject them. It includes a built-in play
28      function that opens the video in your default browser at the exact rally segment (using HTML5
29      media fragments `#t=start,end` for local files or YouTube start time markers `&t=start` for
30      remote videos) for manual visual verification.
31
32    ---
33
34    ## Installation
35
36    Ensure you have Python 3.9+ installed.
37
38    1. Clone or download the repository.
39    2. Install the package and dependencies:
40    ```bash
41    pip install -e .          # core: curate / import-local / export / status
42    pip install -e ".[fetch]" # adds yt-dlp for video acquisition (fetch / crawler)
43    pip install -e ".[dev]"   # test tooling (pytest, pytest-cov)
44    ```
45
46    ### External requirements (not pip-installable)
47
48    The `fetch` command and the ShuttleSet crawler shell out to external binaries.
49    They degrade gracefully (a clear error / skipped download) when missing, but are
50    required for video acquisition:
51
52    - FFmpeg / ffprobe ? used to probe the real resolution/fps of downloaded files.
53    - A JS runtime (e.g. [deno](https://deno.com)) ? needed by yt-dlp to solve
54      YouTube's challenge for high-resolution formats. `fetch` auto-detects a deno
55      install at `~/deno/bin`.
56
57    Core curation/import/export/status commands work without any of these.
58
59    ---
60
61    ## Usage
62
63    All commands should be executed from the root of the project. Make sure `PYTHONPATH` is
64    set to the current directory:
65
66    ```bash
67    # On Windows PowerShell:
68    $env:PYTHONPATH="."
```

```

54     ``
55
56     ### 1. View Database Status
57     To check the current statistics of staging vs. curated items:
58     ```bash
59     python -m src.cli.curate status
60     ``
61
62     ### 2. Ingest External Datasets to Staging
63     To parse external raw stroke-level datasets (e.g. CoachAI format) and register them in
64     staging:
65     ```bash
66     python -m src.cli.ingest --sport badminton --parser coachai --file-path
67     data/staging_coachai_match.csv --license CC-BY-NC-SA --video-id staging_match_non_comm
68     ``
69
70     ### 3. Launch the Interactive Curation Tool & Play Harness
71     To review staging items, inspect their timestamps, play individual rallies in the
72     browser to visually verify their start and end bounds, and approve or reject them:
73     ```bash
74     python -m src.cli.curate curate
75     ``
76
77     ### 4. Manually Import Local Videos & Annotations
78     To register a local video file alongside its JSON/CSV annotation file (bypasses
79     staging):
80     ```bash
81     python -m src.cli.curate import-local --sport badminton --video-path
82     "data/videos/my_custom_match.mp4" --annotations-path "data/mock_local_badminton_anno.json"
83     --match-name "My Custom Local Match" --license Proprietary
84     ``
85
86     `import-local` works for every configured sport ? `badminton`, `tennis`, `cricket`,
87     `pickleball`, and `tabletennis` ? from either a JSON or a CSV annotations file. For the
88     canonical rally CSV schema and the rater-facing instructions used to produce these files (e.g.
89     via an annotation vendor), see [docs/annotation/](docs/annotation/).
90
91     ### 5. Fetch / Upgrade Videos to Best Resolution
92     Videos are always downloaded at the best resolution available (via yt-dlp `bv*+ba/b
93     -S res,...`, merged to MP4); the actual width/height/fps is ffprobed and stored ? never assumed.
94     The `fetch` command (re)downloads videos for entries already in the DB and updates
95     `local_path`/`resolution`/`fps` in place, preserving all rally annotations and curation
96     status (no clean slate needed):
97     ```bash
98     # Upgrade every sub-720p video to best resolution, using a logged-in browser's cookies:
99     python -m src.cli.curate fetch --cookies-from-browser firefox
100
101     # A single video, forcing a re-download even if a good local copy exists:
102     python -m src.cli.curate fetch --video-id shuttleset_1 --force --cookies-from-browser
103     firefox
104     ``
105
106     Flags: `--sport`, `--video-id`, `--status {staging,curated,all}` (default `all`),
107     `--min-height` (skip videos already at/above it unless `--force`; default `720`), `--max-height`
108     (cap resolution, e.g. `1080`; default best available), `--force`, `--include-noncommercial`,
109     `--cookies-from-browser`, `--cookies`, `--player-client`. A global resolution cap can also be set
110     via `video_storage.max_height` in `config/config.yaml`. Re-runs are cheap and resumable; a
111     download that falls below `--min-height` is treated as a failure (discarded, DB left
112     unchanged) so a flaky low-res result never overwrites a good entry.
113
114     > Authenticated cookies are strongly recommended. Anonymous YouTube downloads are
115     throttled and non-deterministic (often falling back to 144p). Pass `--cookies-from-browser
116     firefox` (or `chrome`/`edge`/`brave`, fully closed) ? or `--cookies cookies.txt` ? to use a
117     logged-in (ideally Premium) session for fast, reliable best-resolution downloads. Chrome's app-
118     bound cookie encryption may block direct reads on Windows; Firefox reads cleanly even while
119     open.
120
121     >

```

```

94     > **JS runtime:** yt-dlp needs one (e.g. [deno](https://deno.com)) to solve YouTube's
challenges for high-res formats; `fetch` auto-detects a deno install at `~/deno/bin`. The
default `--player-client` avoids `web`/`web_safari`, which YouTube forces into SABR streaming
(drops to 144p).
95
96     ### 6. Export Annotations as Indexer-Ready Eval/Training Sets
97     To turn curated annotations into ground-truth files the badminton-highlight-indexer can
score and train against, export per-video `start,end` CSVs plus a `manifest.json`:
98     ```bash
99     # Curated + commercially-safe videos across all supported sports (default):
100     python -m src.cli.curate export --out-dir data/eval_export
101
102     # One sport, also including staging entries:
103     python -m src.cli.curate export --sport badminton --include-staging
104     ```
105     The exporter writes `//.csv` (decimal-second `start,end` rows)
and a `manifest.json` pairing each CSV with its video `local_path`/`video_url`, fps, resolution,
and license. The indexer keys videos by **file MD5**, not by our `video_id`, so the downstream
eval/training run must process the exact video file named in the manifest. Flags: `--sport`
(restrict to one sport), `--include-staging` (default: curated only), `--include-noncommercial`
(default: commercial-safe only).
106
107     ---
108
109     ## Configuration
110
111     Adjust settings in `config/config.yaml`:
112     - **`database_path`**: Path to the SQLite database.
113     - **`video_storage`**: Local folder path and automatic download size limit.
114     - **`commercial_licenses`**: Whitelist of licenses allowed for commercial use.
115     - **`supported_sports`**: Allowed sports categories.
116
117     ---
118
119     ## Documentation
120
121     - [docs/PIPELINE_ROLE.md](docs/PIPELINE_ROLE.md) ? how this repo feeds the **badminton-
highlight-indexer** (`export` / `fetch`).
122     - [docs/DECISIONS.md](docs/DECISIONS.md) ? collector-local decision records
(**ADR-001**: keep explicit per-sport code instead of a registry, with an "adding a new sport"
checklist).
123     - [docs/annotation/](docs/annotation/) ? the vendor annotation package used to produce
golden rally labels:
124         - [R00RA_BRIEF.md](docs/annotation/R00RA_BRIEF.md) ? vendor brief + the canonical
rally CSV schema.
125         - [ANNOTATION_GUIDE.md](docs/annotation/ANNOTATION_GUIDE.md) ? the illustrated rater
guide (singles & doubles).
126         - [INTAKE_MANIFEST_TEMPLATE.csv](docs/annotation/INTAKE_MANIFEST_TEMPLATE.csv) ? per-
video intake sheet.
127         - [NDA_TEMPLATE.md](docs/annotation/NDA_TEMPLATE.md) ? mutual NDA template (have
counsel review).
128
129     ---
130
131     ## Running Tests
132
133     We use `pytest` for unit and integration testing:
134     ```bash
135     python -m pytest tests/
136     python -m pytest tests/ --cov=src --cov-report=term-missing # with coverage (~87%)
137     ```
138

```

No files found

8. user (2026-07-02T04:46:05.684Z)

On branch master

Your branch is up to date with 'origin/master'.

Changes not staged for commit:

(use "git add <file>..." to update what will be committed)

(use "git restore <file>..." to discard changes in working directory)

modified: .gitignore

no changes added to commit (use "git add" and/or "git commit -a")

adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)

6f67d65 feat(import): callable import_annotations + a maturity gate (P9) (#52)

c281f82 feat(export): add source file MD5 to the eval/training manifest (P9) (#51)

a152cf5 fix(cvat): warn on rows silently skipped by load_canonical_csv (#50)

75bb0fa docs: archive CODE_HEALTH_PLAN ? project DONE (#49)

origin https://github.com/Khelsutra/sports-data-collector.git (fetch)

origin https://github.com/Khelsutra/sports-data-collector.git (push)

9. user (2026-07-02T04:46:17.670Z)

Already up to date.

adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)

10. user (2026-07-02T04:46:26.031Z)

Preparing worktree (new branch 'fix/curate-verify-registry-rot')

branch 'fix/curate-verify-registry-rot' set up to track 'origin/master'.

HEAD is now at adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)

11. user (2026-07-02T04:46:33.835Z)

C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/config:
config.yaml

C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/data:
mock_local_badminton_anno.json
roora_badminton_canonical.csv
roora_badminton_issues.txt
sports_metadata.db
staging_coachai_match.csv
validation

C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/src/cli:
curate.py
ingest.py
normalize_annotations.py

C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/src/core:
batch_preflight.py
config.py
db.py
downloader.py
manifests.py
parsing.py
proxy.py
schemas

validator.py

C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/tests:
test_batch_preflight.py
test_cleanup.py
test_cli.py
test_config.py
test_crawler_tempfile.py
test_cvat_parser.py
test_cvat_robustness.py
test_db.py
test_db_pragmas.py
test_downloader.py
test_import_guard.py
test_license_consistency.py
test_licensing_gate.py
test_local_import.py
test_manifests.py
test_normalize_annotations.py
test_parser.py
test_parsing.py
test_proxy.py
test_reporting.py
test_reporting_faq.py
test_reporting_noraise.py
test_schemas.py
test_validator.py
data/mock_local_badminton_anno.json
data/roora_badminton_canonical.csv
data/roora_badminton_issues.txt
data/sports_metadata.db
data/staging_coachai_match.csv
data/validation/pilot_badminton_01_report.json
data/validation/pilot_badminton_01_report.md
data/validation/pilot_badminton_01_review.csv

12. user (2026-07-02T04:46:44.847Z)

```
columns: ['video_id', 'sport', 'video_url', 'local_path', 'license_type', 'license_url',
'is_commercial', 'status', 'match_name', 'fps', 'resolution', 'annotation_license']
{'video_id': 'adarsh_avi_singles', 'sport': 'badminton', 'video_url': None, 'local_path':
'C:\\Users\\avidu\\Projects\\Annotation Setup\\Collect\\Golden Labelled\\Adarsh and Avi.MP4',
'license_type': 'Proprietary', 'license_url': None, 'is_commercial': 1, 'status': 'curated',
'match_name': 'Adarsh and Avi (singles)', 'fps': 59.94, 'resolution': '1080p',
'annotation_license': None}
{'video_id': 'gbaaddy', 'sport': 'badminton', 'video_url': None, 'local_path':
'C:\\Users\\avidu\\Projects\\Annotation Setup\\Collect\\Golden Labelled\\gBaaddy20May2026.MP4',
'license_type': 'Proprietary', 'license_url': None, 'is_commercial': 1, 'status': 'curated',
'match_name': 'gBaaddy 20 May 2026 (multi-court)', 'fps': 59.94, 'resolution': '1080p',
'annotation_license': None}
{'video_id': 'kushagra_singles', 'sport': 'badminton', 'video_url': None, 'local_path':
'C:\\Users\\avidu\\Projects\\Annotation Setup\\Collect\\Golden
Labelled\\KushagraYashAviFirstVideo.MP4', 'license_type': 'Proprietary', 'license_url': None,
'is_commercial': 1, 'status': 'curated', 'match_name': 'Kushagra Yash Avi (singles, multi-
court)', 'fps': 59.94, 'resolution': '4K', 'annotation_license': None}
{'video_id': 'largetest_doubles', 'sport': 'badminton', 'video_url': None, 'local_path':
'C:\\Users\\avidu\\Projects\\Annotation Setup\\Collect\\Golden Labelled\\LargeTestVideo.MP4',
'license_type': 'Proprietary', 'license_url': None, 'is_commercial': 1, 'status': 'curated',
'match_name': 'LargeTestVideo (doubles)', 'fps': 29.97, 'resolution': '5.3K',
'annotation_license': None}
{'video_id': 'mahadevpura_singles', 'sport': 'badminton', 'video_url': None, 'local_path':
'C:\\Users\\avidu\\Projects\\Annotation Setup\\Collect\\Golden
Labelled\\MahadevpuraSingles.MP4', 'license_type': 'Proprietary', 'license_url': None,
'is_commercial': 1, 'status': 'curated', 'match_name': 'Mahadevpura Singles (borderline multi-
```



```
court)', 'fps': 29.97, 'resolution': '1080p', 'annotation_license': None}
{'video_id': 'my_custom_local_match', 'sport': 'badminton', 'video_url': None, 'local_path':
'data/videos/my_custom_match.mp4', 'license_type': 'Proprietary', 'license_url': None,
'is_commercial': 1, 'status': 'curated', 'match_name': 'My Custom Local Match', 'fps': 30.0,
'resolution': '1080p', 'annotation_license': None}
{'video_id': 'shuttleaset_1', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=0669aZhH0LI', 'local_path': './data/videos\\shuttleaset_1.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Kento MOMOTA CHOU Tien Chen Fuzhou Open 2019 Finals', 'fps': 25.0, 'resolution':
'1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_10', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=li1sbr6S34g', 'local_path': './data/videos\\shuttleaset_10.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'CHOU Tien Chen NG Ka Long Angus Sudirman Cup 2019 Group Stage', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_11', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=yD6WKVqsAKc', 'local_path': './data/videos\\shuttleaset_11.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'CHOU Tien Chen Jonatan CHRISTIE Indonesia Open 2019 Quarter-finals', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_12', 'sport': 'badminton', 'video_url':
'https://youtube.com/watch?v=example', 'local_path': None, 'license_type': 'Unknown',
'license_url': None, 'is_commercial': 0, 'status': 'staging', 'match_name': 'NG Ka Long Angus
SHI Yu Qi Thailand Masters 2020 SemiFinals', 'fps': 30.0, 'resolution': '1080p',
'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_13', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=boQC4J4E1ZQ', 'local_path': './data/videos\\shuttleaset_13.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Kento MOMOTA Viktor AXELSEN Malaysia Masters 2020 Finals', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_14', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=5W6txLGZ1Rs', 'local_path': './data/videos\\shuttleaset_14.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Anders ANTONSEN Jonatan CHRISTIE Indonesia Masters 2020 QuarterFinals', 'fps':
25.0, 'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_15', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=yu9oyMXRGHY', 'local_path': './data/videos\\shuttleaset_15.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Anthony Sinisuka GINTING Anders ANTONSEN Indonesia Masters 2020 Final', 'fps':
25.0, 'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_16', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=5kS_a7vS5xI', 'local_path': './data/videos\\shuttleaset_16.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Anthony Sinisuka GINTING Viktor AXELSEN Indonesia Masters 2020 SemiFinals',
'fps': 25.0, 'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_17', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=G9r400zkz8', 'local_path': './data/videos\\shuttleaset_17.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'NG Ka Long Angus Jonatan CHRISTIE Malaysia Masters 2020 QuarterFinals', 'fps':
25.0, 'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_18', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=8lHAsyRhYYQ', 'local_path': './data/videos\\shuttleaset_18.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Viktor AXELSEN SHI Yu Qi All England Open 2020 QuarterFinals', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_19', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=YmW83aQFAdg', 'local_path': './data/videos\\shuttleaset_19.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Viktor AXELSEN CHEN Long Malaysia Masters 2020 QuarterFinals', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_2', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=-a0I9-Jx0wC', 'local_path': './data/videos\\shuttleaset_2.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'CHEN Long CHOU Tien Chen World Tour Finals Group Stage', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
```

```

{'video_id': 'shuttleaset_20', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=4e3JJ4rvT3Q', 'local_path': './data/videos\\shuttleaset_20.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Viktor AXELSEN NG Ka Long Angus Malaysia Masters 2020 SemiFinals', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_21', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=gloiZ_gTJaE', 'local_path': './data/videos\\shuttleaset_21.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'curated',
'match_name': 'An Se Young Ratchanok Intanon YONEX Thailand Open 2021 QuarterFinals', 'fps':
30.0, 'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_22', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=8u_UHCnYSkK', 'local_path': './data/videos\\shuttleaset_22.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Mia Blichfeldt Busanan Ongbamrungphan YONEX Thailand Open 2021 QuarterFinals',
'fps': 30.0, 'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_23', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=rSK9Qx8LapE', 'local_path': './data/videos\\shuttleaset_23.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Ng Ka Long Angus Lee Cheuk Yiu YONEX Thailand Open 2021 QuarterFinals', 'fps':
30.0, 'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_24', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=NlDrJyQUTSI', 'local_path': './data/videos\\shuttleaset_24.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Anthony Sinisuka Ginting Rasmus Gemke YONEX Thailand Open 2021 QuarterFinals',
'fps': 30.0, 'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_25', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=FZPrpoGdyHI', 'local_path': './data/videos\\shuttleaset_25.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Carolina Marin Supanida Katethong YONEX Thailand Open 2021 QuarterFinals', 'fps':
30.0, 'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_3', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=eugfCRwSBJo', 'local_path': './data/videos\\shuttleaset_3.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Kento MOMOTA CHOU Tien Chen KOREA OPEN 2019 Final', 'fps': 25.0, 'resolution':
'1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_4', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=y6QbtrTV-K0', 'local_path': './data/videos\\shuttleaset_4.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'CHEN Long CHOU Tien Chen Denmark Open 2019 QuarterFinal', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_6', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=SYyvDrUgClc', 'local_path': './data/videos\\shuttleaset_6.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'Kento MOMOTA CHOU Tien Chen Denmark Open 2018 Finals', 'fps': 25.0, 'resolution':
'1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_8', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=32j2Tg64Zbg', 'local_path': './data/videos\\shuttleaset_8.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'CHOU Tien Chen Anders ANTONSEN Fuzhou Open 2019 Semi-finals', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'shuttleaset_9', 'sport': 'badminton', 'video_url':
'https://www.youtube.com/watch?v=8E98Gpk-fOM', 'local_path': './data/videos\\shuttleaset_9.mp4',
'license_type': 'Unknown', 'license_url': None, 'is_commercial': 0, 'status': 'staging',
'match_name': 'CHOU Tien Chen Jonatan CHRISTIE Sudirman Cup 2019 Quarter-finals', 'fps': 25.0,
'resolution': '1080p', 'annotation_license': 'MIT'}
{'video_id': 'staging_match_non_comm', 'sport': 'badminton', 'video_url':
'https://youtube.com/watch?v=example', 'local_path': None, 'license_type': 'CC-BY-NC-SA',
'license_url': None, 'is_commercial': 0, 'status': 'staging', 'match_name':
'staging_coachai_match', 'fps': 30.0, 'resolution': '1080p', 'annotation_license': None}
{'video_id': 'testlarge_short', 'sport': 'badminton', 'video_url': None, 'local_path':
'C:\\Users\\avidu\\Projects\\Annotation Setup\\Collect\\Golden Labelled\\TestLargeVideo.MP4',
'license_type': 'Proprietary', 'license_url': None, 'is_commercial': 1, 'status': 'curated',
'match_name': 'TestLargeVideo (short)', 'fps': 29.97, 'resolution': '5.3K',
'annotation_license': None}

```

13. assistant (2026-07-02T04:46:56.813Z)

All 6 dangling rows and the 2 placeholder rows confirmed. Now verify where the MP4 files actually exist on disk.

14. user (2026-07-02T04:46:58.529Z)

```
Exit code 2
Archive
Golden Labelled
Trajectories
golden_labels_badminton.csv
rally_review.csv
---
Adarsh and Avi.rallies.csv
KushagraYashAviFirstVideo.rallies.csv
LargeTestVideo.rallies.csv
MahadevpuraSingles.rallies.csv
TestLargeVideo.rallies.csv
gBaaddy20May2026.rallies.csv
---
ls: cannot access 'C:/Users/avidu/Projects/khelsutra-guru/Annotation Setup/Collect/': No such
file or directory
```

15. user (2026-07-02T04:46:59.397Z)

```
ls: cannot access '/f/': No such file or directory
---
$Recycle.Bin
$SysReset
Documents and Settings
DumpStack.log.tmp
F3_Upgrade
```

16. user (2026-07-02T04:47:07.743Z)

Name	Root	Used
Free		
----	----	----

C	C:\	534028824576 266702819328
G	G:\	547363966976 253367676928
Temp	C:\Users\avidu\AppData\Local\Temp\	534028824576 266702819328

17. user (2026-07-02T04:47:07.882Z)

```
123:| original-6 labels | `F:\[Khelsutra]?Training\Golden Labelled\*.rallies.csv` | ~12 KB | ?
| train + eval |
```

18. user (2026-07-02T04:47:14.094Z)

```
123:| original-6 labels | `F:\[Khelsutra]?Training\Golden Labelled\*.rallies.csv` | ~12 KB | ?
| train + eval |
204:> `[Done] Video N/` (Drive) or `[Khelsutra] ?` (F:) fail with "matched no objects." The
owner-box
```

19. user (2026-07-02T04:47:14.821Z)

```
100      --video-uri gcs://khelsutra-rally-corpus/videos/<clip>.mp4 \
101      --out-uri   gcs://khelsutra-rally-corpus --segmenter wasb_hybrid --config config.json
102
103      # TRAIN from the bucket (reads labels/ [+ videos/], extracts trajectories, builds
manifest on-box,
104      # runs training.gen0.harness, pushes weights/<version>/):
```

```

105     RALLY_DETECTOR_IMPL=native python -m backend.worker train \
106         --corpus-uri gcs://khelsutra-rally-corpus --out-uri gcs://khelsutra-rally-corpus \
107         --version <v> --trajectory-source detector --segmenter wasb_hybrid --config
config.json
108     ```
109
110     A bare/relative path resolves under a configured `input_root`/`input_backend`
111     (`backend/storage/ref.py::resolve_input_uri`), so workflows can be written bucket-
relative.
112
113     ---
114
115     ## 5. The gap ? what's still local-only (why the box is still a blocker)
116
117     All of this is tiny except the originals ? the metadata that blocks cloud eval is <
25 MB total:
118
119     | Data | Where it is now | Size | In bucket? | Needed for |
120     |---|---|---|---|---|
121     | 20 WASB trajectory CSVs | `C:\?\Annotation Setup\Collect\Trajectories\` | 17.5 MB
| ? none | eval/litmus (cached). *(No `cached` train source exists ? train synths or detector-
extracts.)* |
122     | golden manifest | `output/human_lovo_manifest.json` (C:-absolute paths) | 4.5 KB | ? |
eval corpus definition |
123     | original-6 labels | `F:[Khelsutra]?\Training\Golden Labelled\*.rallies.csv` | ~12 KB
| ? | train + eval |
124     | 7 new-clip labels | repo `output/*.rallies.csv` (gitignored) F: archive (backed
up 2026-06-22) | ~12 KB | ? partial/inconsistent | train + eval |
125     | source videos / proxies | F: archive + Google Drive golden folder | 100s of GB | ?
3 only | infer + train (real WASB extraction) |
126
127     > The corpus metadata (trajectories + labels + manifest, ~25 MB) is what actually
blocks GPU-free
128     > cloud eval and the served Gate-A litmus ? and it's trivially small to migrate. The
videos are the
129     > only large migration, and they already exist on Google Drive (the golden-corpus
folder), so they
130     > can go Drive ? bucket from a cloud box, keeping the local machine out of the path
entirely.
131
132     ---
133
134     ## 6. Migration plan (close the gap)
135
136     Phase 1 ? corpus metadata (?25 MB, do first, removes the eval blocker). Stage all
trajectories,
137     all labels, and a bucket-relative eval manifest into `gs://khelsutra-rally-corpus`.
Pending the §7
138     naming decision. Sketch:
139
140     ```bash
141     # trajectories (decision-free; absent today)
142     gcloud storage cp "C:\?\Collect\Trajectories\*.csv" gs://khelsutra-rally-
corpus/trajectories/
143     # labels (after the §7.1 naming decision) -> labels/<canonical-name>.rallies.csv
144     # eval manifest with gs:// (or bucket-relative) paths -> a stable manifest object (§7.3)
145     ```
146
147     Phase 2 ? videos ? `videos/<name>.mp4` (bucket name = the eval `name`, consistent
with
148     `trajectories/` + `labels/`). Sources (resolved 2026-06-22 ? see
`scratch/copy_videos_to_gcs.ps1`):
149
150     | Source | Clips | Size |
151     |---|---|---|

```

```

152 | Drive golden folder `~/Golden Label Videos Request - June 15/[Done] Video N/`
(**proxies**) | `mahadevpura_1/2`, `GX010128`, `GX020094`, `GX030094`, `Badminton_BXH_2`,
`Boxhill_Doubles` (7) | ~13.5 GB |
153 | F: archive (**full MP4s** ? proxy-first before upload) | `GX010137`, `GX010141`, +
original-6 (`adarsh_avi_singles`, `kushagra_singles`, `largetest_doubles`, `gbaaddy`,
`testlarge_short`, `mahadevpura_singles`) (8) | ~50 GB |
154
155 **Two ways to run it:**
156 - **Decoupled (preferred ? no local box):** `rclone` on a cloud VM with a Google-Drive
remote ?
157 `rclone copy gdrive:"<?/[Done] Video N>/<proxy>.mp4" :gcs:khelsutra-rally-
corpus/videos/<name>.mp4`
158 (one-time owner step: `rclone config` ? `gdrive` remote OAuth). Or `gdown <file-id>`
on the box for
159 the shared folder. This is the path that actually removes the local box.
160 - **Interim (today, owner box ? streams Drive/F: ? local ? GCS):** `pwsh
scratch/copy_videos_to_gcs.ps1
161 -Execute` (dry-run without `-Execute`; `-Only <name>` for one; skip-existing). Uses
local bandwidth.
162
163 **Priority:** the 7 Drive **proxies** first (already 1080p, sufficient for detection).
The 8 large F:
164 full MP4s should be **proxied to 1080p first** (cheaper storage + faster cloud decode)
before upload.
165 ? Two pre-existing non-canonical video objects (`videos/GX030094_proxy.mp4`,
`videos/mahadevpura-1.MP4`)
166 are superseded once `videos/GX030094.mp4` / `videos/mahadevpura_1.mp4` land ? verify-
dupe-then-delete
167 (same as the label cleanup). Track per-clip in $8.
168
169 **Phase 3 ? make eval/litmus read GCS (real code, NOT a config flip).** The eval drivers
170 (`served_gate_a.evaluate_video`, `calibrate_wasb.load_golden_gts`,
`TrackNetRunner.parse_trajectory_csv`)
171 open `spec["trajectory"]`/`spec["labels"]` with plain `open()` and gate on
`os.path.isfile()` ? they
172 have **no storage-layer fetch**, so a `gs://` manifest cannot be consumed today. Phase-3
must thread
173 `storage.fetch`/`resolve_input_uri` through those paths first. Then: re-probe
`fps`/`frame_width` from
174 the bucket videos (don't trust the copied values ? a wrong fps silently mis-aligns every
label), and
175 **re-measure the served Gate-A litmus on the GCS 15-clip set as a [[launch-report-
convention]] event**
176 (dual-set RAW+golden no-regression, not a silent eval-set swap ? see PR #236).
177
178 ---
179

```

20. assistant (2026-07-02T04:47:31.096Z)

F: is not mounted (only C:, G:, Temp). Let me check the exact documented F: paths in the copy script, and do a thorough search for the six MP4s on the reachable drives (C: workspace and G: Drive).

21. user (2026-07-02T04:47:32.522Z)

```

20 # --- Drive golden proxies (preferred) ---
21 'mahadevpura_2' = "$DRIVE\[Done] Video 3 - Badminton
Multicourt\mahadevpura-2_proxy.mp4"
22 'GX020094' = "$DRIVE\[Done] Video 4 - Badminton\GX020094_proxy.mp4"
23 'GX030094' = "$DRIVE\[Done] Video 6 - Badminton\GX030094_proxy.mp4"
24 'Badminton_BXH_2' = "$DRIVE\[Done] Video 10 - Badminton Multicourt\Badminton BXH
2_proxy.mp4"
25 'Boxhill_Doubles' = "$DRIVE\[Done] Video 12 - Badminton Multicourt\Boxhill
Doubles_proxy.mp4"

```

```

26      'GX010128'          = "$DRIVE\[Done] Video 13 - Badminton\GX010128.MP4"
27      'mahadevpura_1'     = "$DRIVE\[Done] Video 14 - Badminton\mahadevpura-1.MP4"
28      # --- F: full MP4s (large; consider proxy-first before upload) ---
29      'GX010137'          = "$FR00T\Badminton - Boxhill - 17 June\GX010137.MP4"          #
*6.4GB
30      'GX010141'          = "$FR00T\Badminton - Boxhill - 17 June\GX010141.MP4"          #
*5.5GB
31      'adarsh_avi_singles' = "$FR00T\Training\Golden Labelled\Adarsh and Avi.MP4"
# *10.9GB
32      'kushagra_singles'   = "$FR00T\Training\Golden
Labelled\KushagraYashAviFirstVideo.MP4" # *7.7GB
33      'largetest_doubles'  = "$FR00T\Training\Golden Labelled\LargeTestVideo.MP4"
# *11GB
34      'gbaaddy'            = "$FR00T\Training\Golden Labelled\gBaaddy20May2026.MP4"
# *6.7GB
35      'testlarge_short'    = "$FR00T\Training\Golden Labelled\TestLargeVideo.MP4"
# *1.4GB
36      'mahadevpura_singles' = "$FR00T\Training\Golden Labelled\MahadevpuraSingles.MP4"
# 187MB
37      }
38
39      # which dests already exist in the bucket (skip those)
40      $have = @{}
41      gcloud storage ls "$BUCKET/" 2>$null | ForEach-Object {
$have[[IO.Path]::GetFileName($_)] = $true }
42
43      $copied = 0; $skipped = 0; $missing = 0
44      foreach ($name in $src.Keys) {
45          if ($Only -and $name -ne $Only) { continue }
46          $s = $src[$name]; $dst = "$BUCKET/$name.mp4"; $destName = "$name.mp4"
47          $exists = Test-Path -LiteralPath $s
48          $sizeMB = if ($exists) { [math]::Round((Get-Item -LiteralPath $s).Length/1MB) } else {
0 }
49          if (-not $exists) { "MISSING SRC {0,-20} {1}" -f $name, $s; $missing++; continue }

```

22. user (2026-07-02T04:47:46.454Z)

G:\My Drive\Sharing\Badminton\Boxhill\June 1 2026\Adarsh and Avi.MP4 | 11,445,052,772 bytes
G:\My Drive\Sharing\Golden Annotations\MahadevpuraSingles.MP4 | 196,182,380 bytes
search done

23. user (2026-07-02T04:48:55.302Z)

<system-reminder>[Truncated: PARTIAL view ? showing lines 1-1167 of 1532 total (27874 tokens, cap 25000). Call Read with offset=1168 limit=1167 for the next page, or Grep to find a specific section. Do NOT answer from this page alone if the answer may be further in the file.]</system-reminder>

```

1      import os
2      import sys
3      import time
4      import json
5      import csv
6      import argparse
7      from dataclasses import dataclass, field
8      from typing import List, Dict, Any, Optional, Tuple, Union
9      from tabulate import tabulate
10
11      from src.core.config import load_config
12      from src.core.parsing import safe_int as _safe_int, safe_float_required as _safe_float
13      from src.core.db import SportsDatabase
14      from src.core.manifests import EvalSetManifest
15      from src.core.schemas.base import VideoMetadata, SportType, LicenseType, CurationStatus
16      from src.core.schemas.badminton import BadmintonRally
17      from src.core.schemas.tennis import TennisPoint

```

```

18     from src.core.schemas.cricket import CricketDelivery
19     from src.core.schemas.pickleball import PickleballRally
20     from src.core.schemas.tabletennis import TableTennisRally
21     from src.core.validator import DatasetValidator
22     from src.core.downloader import download_best_video, probe_resolution_fps,
height_to_label
23
24     #: A parsed annotation row of any sport ? the union of all segment schema types.
25     #: Used so _parse_local_annotations / _guard_import return concrete types, not Any.
26     Annotation = Union[BadmintonRally, TennisPoint, CricketDelivery,
27                        PickleballRally, TableTennisRally]
28
29
30     class CurationError(Exception):
31         """A user-correctable pre-flight failure (bad input, duplicate ids, clobber guard).
32
33         Raised by library-style methods (e.g. ``_guard_import``) instead of calling
34         ``sys.exit(1)`` directly, so the logic is reusable from tests or a future API.
35         The CLI ``main()`` translates it into ``sys.exit(1)``; a library caller can
36         catch and handle it. See docs/CODE_HEALTH_PLAN.md §7 Tier D (A2).
37         """
38
39
40     @dataclass
41     class ImportResult:
42         """The structured outcome of ``import_annotations`` ? so a service/UI caller gets
data back
43         instead of stdout + sys.exit. ``status`` is the CurationStatus the video was
recorded with (the
44         maturity gate: ``staging`` = needs review, ``curated`` = confirmed)."""
45         video_id: str
46         count: int
47         status: str
48         match_name: str
49         is_commercial: bool
50         license_type: str
51         warnings: List[str] = field(default_factory=list)
52         video_meta: Optional[VideoMetadata] = None
53         annotations: List[Annotation] = field(default_factory=list)
54
55
56     def _truthy(value: Any) -> bool:
57         """Coerce a CSV/JSON cell to a boolean, treating "0"/"false"/"no"/"" as False.
58
59         B4: ``bool("0")`` is True (non-empty string), so the JSON path's bare ``bool(...)``
60         disagreed with the CSV path's explicit truthiness list on the same logical value.
61         Both paths now route through this single helper.
62         """
63         if isinstance(value, bool):
64             return value
65         if value is None:
66             return False
67         return str(value).strip().lower() in ("true", "1", "yes")
68
69
70     class SportsDataCLI:
71         def __init__(self, config_path: str = "config/config.yaml"):
72             self.config_path = config_path
73             self.config = self._load_config()
74
75             # Resolve DB path
76             db_path = self.config.get("database_path", "data/sports_metadata.db")
77             self.db = SportsDatabase(db_path)
78             self.validator = DatasetValidator(config_path)
79

```

```

80     def _load_config(self) -> Dict[str, Any]:
81         if not os.path.exists(self.config_path):
82             print(f"Warning: Configuration file not found at {self.config_path}. Using
defaults.")
83             return load_config(self.config_path)
84
85     def status(self):
86         """Displays status summary of all videos in the database."""
87         # Query total count by status via the public DB read API.
88         rows = self.db.get_status_counts() # list of (sport, status, cnt)
89
90         if not rows:
91             print("\nDatabase is currently empty. Run crawlers or import local data.\n")
92             return
93
94         data = [[sport, status, cnt] for sport, status, cnt in rows]
95         print("\n=== Dataset Ingestion Status ===")
96         print(tabulate(data, headers=["Sport", "Status", "Count"], tablefmt="grid"))
97         print()
98
99     @staticmethod
100     def _annotation_key(sport_str: str, a: Annotation) -> Union[int, Tuple[int, int],
None]:
101         """The per-sport primary-key value of one annotation (for duplicate
detection)."""
102         if sport_str in ("badminton", "pickleball", "tabletennis"):
103             return getattr(a, "rally_number", None)
104         if sport_str == "tennis":
105             return getattr(a, "point_number", None)
106         if sport_str == "cricket":
107             return (getattr(a, "over", None), getattr(a, "ball", None))
108         return None
109
110     def _guard_import(self, sport_str: str, video_id: str, annotations:
List[Annotation],
111                      force: bool = False) -> None:
112         """Pre-insert safety checks (run BEFORE any DB write, so a rejection is
atomic)."""
113         # (1) Duplicate primary-key guard. int(float()) truncation (e.g. 7 and 7.5 both
114         # become 7) would otherwise raise IntegrityError mid-insert, leaving a committed
115         # video row with rolled-back annotations. Catch it before touching the DB.
116         keys = [self._annotation_key(sport_str, a) for a in annotations]
117         keys = [k for k in keys if k is not None]
118         seen, dupes = set(), set()
119         for k in keys:
120             if k in seen:
121                 dupes.add(k)
122             seen.add(k)
123         if dupes:
124             raise CurationError(
125                 f"Duplicate rally/point identifier(s) {sorted(dupes)} in the annotations
"
126                 f"(often a fractional id like 7.5 truncating to 7). "
127                 f"Re-run with --normalize, or normalize first: "
128                 f"python -m src.cli.normalize_annotations --in <csv> --out clean.csv")
129
130         # (2) Golden-label protection ? refuse to silently clobber an existing video's
131         # annotations (filename-derived video_ids collide easily).
132         if self.db.get_video(video_id):
133             try:
134                 n = len(self.db.get_segment_intervals(video_id, SportType(sport_str)))
135             except Exception:
136                 n = 0
137             if n > 0 and not force:
138                 raise CurationError(

```



```

139             f"video_id '{video_id}' already has {n} annotated segment(s); "
140             f"importing would REPLACE them. "
141             f"Pass --force to overwrite, or use a distinct --video-id to keep
both.")
142         if n > 0:
143             print(f"Warning: --force given ? replacing {n} existing segment(s) for
video_id '{video_id}'.")
144
145     def import_annotations(
146         self,
147         *,
148         sport: str,
149         video_path: Optional[str],
150         annotations_path: str,
151         video_id: Optional[str] = None,
152         match_name: Optional[str] = None,
153         license: str = "Proprietary",
154         fps: Optional[float] = None,
155         resolution: Optional[str] = None,
156         normalize: bool = False,
157         force: bool = False,
158         status: CurationStatus = CurationStatus.CURATED,
159     ) -> ImportResult:
160         """Ingest one video's rally labels into the DB and RETURN a structured result ?
the callable
161         core behind the CLI ``import-local`` (and a future UI/service contribute path).
Raises
162         ``CurationError`` on any user-correctable failure (unsupported sport,
missing/bad annotations,
163         validation, duplicate/clobber guard); it prints nothing and never calls
``sys.exit``, so a
164         service can catch and surface it.
165
166         ``status`` is the MATURITY GATE: pass ``CurationStatus.STAGING`` for a
contributed label that
167         NEEDS REVIEW (the safe default for programmatic/UI imports ? never auto-promote
to golden),
168         ``CURATED`` once confirmed. The CLI passes CURATED (a manual owner import is
reviewer-vetted).
169         """
170         warnings: List[str] = []
171         sport_str = (sport or "").lower()
172         if sport_str not in self.config.get("supported_sports", ["badminton"]):
173             raise CurationError(
174                 f"Sport '{sport_str}' is not supported. Supported:
{self.config.get('supported_sports')}")
175
176         if video_path and not os.path.exists(video_path):
177             warnings.append(f"Local video path '{video_path}' does not exist on disk,
but continuing registration.")
178         if not annotations_path or not os.path.exists(annotations_path):
179             raise CurationError(f"Annotations file '{annotations_path}' not found.")
180
181         match_name = match_name or os.path.basename(video_path or
annotations_path).split('.')[0]
182         vid = video_id or match_name.lower().replace(" ", "_")
183
184         try:
185             license_type = LicenseType(license)
186         except ValueError:
187             warnings.append(f"Unknown license '{license}'. Defaulting to Unknown (non-
commercial).")
188             license_type = LicenseType.UNKNOWN
189         is_commercial = self.validator.check_license_commercial(license_type)
190

```

```

191         # Optional: normalize a raw CSV annotator export before parsing (drops unusable
rows,
192         # rennumbers fractional/duplicate rally ids) so the guards below can't fire.
193         parse_path = annotations_path
194         if normalize and parse_path.lower().endswith(".csv"):
195             from src.cli.normalize_annotations import normalize_file
196             normalized_path = parse_path + ".normalized.csv"
197             rep = normalize_file(parse_path, normalized_path)
198             warnings.append("normalize-annotations:\n" + rep.summary())
199             parse_path = normalized_path
200
201         try:
202             annotations = self._parse_local_annotations(sport_str, vid, parse_path)
203         except Exception as e:
204             raise CurationError(f"Error parsing annotations: {e}") from e
205
206         is_valid, errors = True, []
207         if sport_str == "badminton":
208             is_valid, errors = self.validator.validate_badminton_rallies(annotations)
209         elif sport_str == "tennis":
210             is_valid, errors = self.validator.validate_tennis_points(annotations)
211         elif sport_str == "cricket":
212             is_valid, errors = self.validator.validate_cricket_deliveries(annotations)
213         elif sport_str == "pickleball":
214             is_valid, errors = self.validator.validate_pickleball_rallies(annotations)
215         elif sport_str == "tabletennis":
216             is_valid, errors = self.validator.validate_tabletennis_rallies(annotations)
217         if not is_valid:
218             raise CurationError("Annotations failed validation:\n" + "\n".join(f" -
{e}" for e in errors))
219
220         # Pre-insert safety (raises CurationError) ? atomic: no half-written DB on
rejection.
221         self._guard_import(sport_str, vid, annotations, force=force)
222
223         video_meta = VideoMetadata(
224             video_id=vid, sport=SportType(sport_str), video_url=None,
local_path=video_path,
225             license_type=license_type, license_url=None, is_commercial=is_commercial,
226             status=status, match_name=match_name, fps=fps, resolution=resolution,
227         )
228         self.db.insert_video(video_meta)
229         count = 0
230         if sport_str == "badminton":
231             count = self.db.insert_badminton_rallies(annotations)
232         elif sport_str == "tennis":
233             count = self.db.insert_tennis_points(annotations)
234         elif sport_str == "cricket":
235             count = self.db.insert_cricket_deliveries(annotations)
236         elif sport_str == "pickleball":
237             count = self.db.insert_pickleball_rallies(annotations)
238         elif sport_str == "tabletennis":
239             count = self.db.insert_tabletennis_rallies(annotations)
240
241         return ImportResult(
242             video_id=vid, count=count, status=status.value, match_name=match_name,
243             is_commercial=is_commercial, license_type=license_type.value,
244             warnings=warnings, video_meta=video_meta, annotations=annotations,
245         )
246
247     def import_local(self, args):
248         """CLI: manually register a local video + annotations (a curated owner import).
Thin wrapper
249         over ``import_annotations`` ? prints the warnings/summary, emits the
observability report, and

```

```

250         translates a ``CurationError`` into ``sys.exit(1)``. """
251     try:
252         res = self.import_annotations(
253             sport=args.sport, video_path=args.video_path,
annotations_path=args.annotations_path,
254             video_id=args.video_id, match_name=args.match_name,
license=args.license,
255             fps=args.fps, resolution=args.resolution,
256             normalize=getattr(args, "normalize", False), force=getattr(args,
"force", False),
257             status=CurationStatus.CURATED, # a manual owner import is already
reviewer-vetted
258         )
259     except CurationError as e:
260         print(f"Error: {e}")
261         sys.exit(1)
262
263     for w in res.warnings:
264         print(f"Warning: {w}")
265     print(f"\nSuccessfully imported local video '{res.match_name}':")
266     print(f"  - Video ID: {res.video_id}")
267     print(f"  - Sport: {args.sport.lower()}")
268     print(f"  - Commercial Use Allowed: {res.is_commercial} (License:
{res.license_type})")
269     print(f"  - Registered {res.count} annotated segments.\n")
270
271     # Per-video observability report (next to the video). Never raises.
272     from src.reporting import emit_import_report
273     src_fmt = os.path.splitext(args.annotations_path)[1].lstrip(".").lower() or
"unknown"
274     report_paths = emit_import_report(
275         video_meta_obj=res.video_meta, annotations=res.annotations, is_valid=True,
276         errors=[], count=res.count, sport=args.sport.lower(), source_format=src_fmt,
277         fallback_dir=os.path.dirname(os.path.abspath(args.annotations_path)),
278     )
279     if report_paths:
280         print(f"  - Observability report: {report_paths.get('technical_md')}")
281         print(f"  - Customer summary      : {report_paths.get('customer_md')}\n")
282
283     def _parse_local_annotations(self, sport: str, video_id: str, filepath: str) ->
List[Annotation]:
284         """Parses simple JSON or CSV annotations files."""
285         parsed_items = []
286         ext = os.path.splitext(filepath)[1].lower()
287
288         if ext == '.json':
289             with open(filepath, 'r', encoding='utf-8') as f:
290                 data = json.load(f)
291
292             if not isinstance(data, list):
293                 raise ValueError("JSON file must contain a list of objects.")
294
295             # Uses the same safe casts as the CSV path: a missing/blank required
296             # time raises a clear ValueError (caught by the caller) instead of a
297             # bare KeyError, and optional ints degrade to sensible defaults.
298             for idx, obj in enumerate(data):
299                 if sport == "badminton":
300                     parsed_items.append(BadmintonRally(
301                         video_id=video_id,
302                         rally_number=_safe_int(obj.get("rally_number"), idx + 1),
303                         start_time=_safe_float(obj.get("start_time")),
304                         end_time=_safe_float(obj.get("end_time")),
305                         score_before=obj.get("score_before"),
306                         score_after=obj.get("score_after"),
307                         shots_count=_safe_int(obj.get("shots_count"), None),

```

```

308         ending_reason=obj.get("ending_reason"),
309         last_shot_outcome=obj.get("last_shot_outcome")
310     ))
311 elif sport == "tennis":
312     parsed_items.append(TennisPoint(
313         video_id=video_id,
314         point_number=_safe_int(obj.get("point_number"), idx + 1),
315         start_time=_safe_float(obj.get("start_time")),
316         end_time=_safe_float(obj.get("end_time")),
317         set_score=obj.get("set_score"),
318         game_score=obj.get("game_score"),
319         server=obj.get("server"),
320         ending_type=obj.get("ending_type")
321     ))
322 elif sport == "cricket":
323     parsed_items.append(CricketDelivery(
324         video_id=video_id,
325         over=_safe_int(obj.get("over"), 0),
326         ball=_safe_int(obj.get("ball"), 0),
327         start_time=_safe_float(obj.get("start_time")),
328         end_time=_safe_float(obj.get("end_time")),
329         batsman=obj.get("batsman"),
330         bowler=obj.get("bowler"),
331         runs=_safe_int(obj.get("runs"), None),
332         # B4: route through the same truthiness helper as the CSV path
333         # below ? bool("0") is True (non-empty string), so the JSON and
334         # CSV paths previously disagreed on the same logical value.
335         wicket=_truthy(obj.get("wicket")),
336         extras=_safe_int(obj.get("extras"), 0)
337     ))
338 elif sport == "pickleball":
339     parsed_items.append(PickleballRally(
340         video_id=video_id,
341         rally_number=_safe_int(obj.get("rally_number"), idx + 1),
342         start_time=_safe_float(obj.get("start_time")),
343         end_time=_safe_float(obj.get("end_time")),
344         score_before=obj.get("score_before"),
345         score_after=obj.get("score_after"),
346         shots_count=_safe_int(obj.get("shots_count"), None),
347         ending_reason=obj.get("ending_reason"),
348         last_shot_outcome=obj.get("last_shot_outcome")
349     ))
350 elif sport == "tabletennis":
351     parsed_items.append(TableTennisRally(
352         video_id=video_id,
353         rally_number=_safe_int(obj.get("rally_number"), idx + 1),
354         start_time=_safe_float(obj.get("start_time")),
355         end_time=_safe_float(obj.get("end_time")),
356         score_before=obj.get("score_before"),
357         score_after=obj.get("score_after"),
358         shots_count=_safe_int(obj.get("shots_count"), None),
359         ending_reason=obj.get("ending_reason"),
360         last_shot_outcome=obj.get("last_shot_outcome")
361     ))
362 elif ext == '.csv':
363     with open(filepath, 'r', encoding='utf-8') as f:
364         reader = csv.DictReader(f)
365         # Normalize headers
366         reader.fieldnames = [n.lower().strip() for n in reader.fieldnames] if
reader.fieldnames else []
367
368     for idx, row in enumerate(reader):
369         if sport == "badminton":
370             parsed_items.append(BadmintonRally(
371                 video_id=video_id,

```

```

372         rally_number=_safe_int(row.get("rally_number"), idx + 1),
373         start_time=_safe_float(row.get("start_time")),
374         end_time=_safe_float(row.get("end_time")),
375         score_before=row.get("score_before"),
376         score_after=row.get("score_after"),
377         shots_count=_safe_int(row.get("shots_count"), None),
378         ending_reason=row.get("ending_reason"),
379         last_shot_outcome=row.get("last_shot_outcome")
380     ))
381     elif sport == "tennis":
382         parsed_items.append(TennisPoint(
383             video_id=video_id,
384             point_number=_safe_int(row.get("point_number"), idx + 1),
385             start_time=_safe_float(row.get("start_time")),
386             end_time=_safe_float(row.get("end_time")),
387             set_score=row.get("set_score"),
388             game_score=row.get("game_score"),
389             server=row.get("server"),
390             ending_type=row.get("ending_type")
391         ))
392     elif sport == "cricket":
393         parsed_items.append(CricketDelivery(
394             video_id=video_id,
395             over=_safe_int(row.get("over"), 0),
396             ball=_safe_int(row.get("ball"), 0),
397             start_time=_safe_float(row.get("start_time")),
398             end_time=_safe_float(row.get("end_time")),
399             batsman=row.get("batsman"),
400             bowler=row.get("bowler"),
401             runs=_safe_int(row.get("runs"), None),
402             wicket=_truthy(row.get("wicket")),
403             extras=_safe_int(row.get("extras"), 0)
404         ))
405     elif sport == "pickleball":
406         parsed_items.append(PickleballRally(
407             video_id=video_id,
408             rally_number=_safe_int(row.get("rally_number"), idx + 1),
409             start_time=_safe_float(row.get("start_time")),
410             end_time=_safe_float(row.get("end_time")),
411             score_before=row.get("score_before"),
412             score_after=row.get("score_after"),
413             shots_count=_safe_int(row.get("shots_count"), None),
414             ending_reason=row.get("ending_reason"),
415             last_shot_outcome=row.get("last_shot_outcome")
416         ))
417     elif sport == "tabletennis":
418         parsed_items.append(TableTennisRally(
419             video_id=video_id,
420             rally_number=_safe_int(row.get("rally_number"), idx + 1),
421             start_time=_safe_float(row.get("start_time")),
422             end_time=_safe_float(row.get("end_time")),
423             score_before=row.get("score_before"),
424             score_after=row.get("score_after"),
425             shots_count=_safe_int(row.get("shots_count"), None),
426             ending_reason=row.get("ending_reason"),
427             last_shot_outcome=row.get("last_shot_outcome")
428         ))
429     else:
430         raise ValueError("Unsupported annotations file extension. Use .json or
431 .csv.")
432     return parsed_items
433
434 def _fetch_one_with_retry(self, url, dest_path, args, max_height, max_size_mb):
435     """Download one video at best resolution, retrying transient failures.

```

```

436
437     A below-min-height result almost always means YouTube rate-limited the
438     request and served a low fallback (not that the video lacks HD), so we
439     retry with a growing backoff. Returns (DownloadResult, None) on success,
440     or (None, failure_reason) once retries are exhausted. A below-min file is
441     discarded each attempt so it never lands in the DB.
442     """
443     attempts = max(1, args.retries + 1)
444     reason = "unknown error"
445     for i in range(attempts):
446         result = download_best_video(
447             url, dest_path,
448             max_height=max_height, max_size_mb=max_size_mb,
449             cookies_from_browser=args.cookies_from_browser,
450             cookies_file=args.cookies,
451             player_client=args.player_client,
452         )
453         if result.success and (result.height is None or result.height >=
args.min_height):
454             return result, None
455
456         if not result.success:
457             reason = result.error or "unknown error"
458         else:
459             reason = (f"got {result.resolution_label} (< min-height
{args.min_height}p) ? "
460                      f"YouTube served a low format (rate-limited?)")
461             try:
462                 os.remove(result.path)
463             except OSError:
464                 pass
465
466             if i < attempts - 1:
467                 backoff = args.retry_backoff * (i + 1)
468                 print(f" attempt {i+1}/{attempts} failed: {reason}\n retrying in
{backoff:.0f}s...")
469                 time.sleep(backoff)
470             return None, reason
471
472     def fetch(self, args):
473         """(Re)download videos at best available resolution, updating the DB in place.
474
475         This is the in-place upgrade path: rows, rally annotations, and curation
476         status are all preserved ? only the video file plus its `local_path`,
477         `resolution`, and `fps` are refreshed. By default it skips videos that
478         already have a local file at or above `--min-height`, so re-runs are cheap
479         and resumable; `--force` re-fetches everything.
480         """
481         storage = self.config.get("video_storage", {})
482         local_dir = storage.get("local_dir", "./data/videos")
483         max_download_size_mb = storage.get("max_download_size_mb")
484         max_height = args.max_height if args.max_height is not None else
storage.get("max_height")
485         os.makedirs(local_dir, exist_ok=True)
486
487         commercial_only = not args.include_noncommercial
488         if args.status == "all":
489             statuses = [CurationStatus.CURATED, CurationStatus.STAGING]
490         else:
491             statuses = [CurationStatus(args.status)]
492
493         supported = self.config.get("supported_sports", ["badminton"])
494         sport_strs = [args.sport] if args.sport else supported
495
496         # Collect target videos.

```

```

497     videos = []
498     for sport_str in sport_strs:
499         try:
500             sport = SportType(sport_str.lower())
501         except ValueError:
502             print(f"Warning: skipping unknown sport '{sport_str}'.")
503             continue
504         for status in statuses:
505             videos.extend(self.db.get_videos_by_status(status, sport))
506     if args.video_id:
507         videos = [v for v in videos if v.video_id == args.video_id]
508
509     if not videos:
510         print("\nNo matching videos to fetch.\n")
511         return
512
513     results = [] # (video_id, action, detail)
514     for video in videos:
515         if commercial_only and not video.is_commercial:
516             results.append((video.video_id, "skip", "non-commercial"))
517             continue
518         if not video.video_url or "watch?v=example" in video.video_url:
519             # Either no URL, or the parser's placeholder (the ShuttleSet catalog
520             # had no URL for this match) ? nothing to download.
521             results.append((video.video_id, "skip", "no real source URL"))
522             continue
523
524         dest_path = os.path.join(local_dir, f"{video.video_id}.mp4")
525
526         # Skip if we already have a good-enough local file (unless --force).
527         if not args.force and os.path.exists(dest_path):
528             _, height, _ = probe_resolution_fps(dest_path)
529             if height is not None and height >= args.min_height:
530                 results.append((video.video_id, "have", f"{height_to_label(height)}
already"))
531                 continue
532
533         print(f"Fetching {video.video_id} ({video.match_name})...")
534         result, fail_reason = self._fetch_one_with_retry(
535             video.video_url, dest_path, args, max_height, max_download_size_mb)
536         if fail_reason is not None:
537             results.append((video.video_id, "FAIL", fail_reason))
538             if args.sleep:
539                 time.sleep(args.sleep)
540             continue
541
542         # Update only the file-derived fields; preserve everything else (UPSERT
543         # mutates the row in place, leaving rallies intact).
544         video.local_path = result.path
545         if result.resolution_label:
546             video.resolution = result.resolution_label
547         if result.fps:
548             video.fps = result.fps
549         self.db.insert_video(video)
550
551         detail = (f"{result.resolution_label or '?'}"
552                 f"{f' {result.fps:g}fps' if result.fps else ''},
{result.size_mb:.0f} MB")
553         results.append((video.video_id, "fetched", detail))
554         if args.sleep:
555             time.sleep(args.sleep)
556
557     print("\n=== Fetch summary ===")
558     print(tabulate([[vid, action, detail] for vid, action, detail in results],
559                   headers=["Video ID", "Action", "Detail"], tablefmt="grid"))

```

```

560         fetched = sum(1 for _, a, _ in results if a == "fetched")
561         failed = [r for r in results if r[1] == "FAIL"]
562         print(f"\nFetched/updated : {fetched}    Already-good : {sum(1 for _,a,_ in
results if a=='have')}}    "
563             f"Skipped : {sum(1 for _,a,_ in results if a=='skip')}}    Failed :
{len(failed)}}")
564         if failed:
565             print("Failures (left unchanged in DB):")
566             for vid, _, detail in failed:
567                 print(f"    - {vid}: {detail}")
568             print()
569
570         # Code/data licenses that can NEVER apply to video footage ? if one of these is
571         # stamped on a remote/broadcast video it is an annotation-vs-video license
conflation.
572         _CODE_DATA_LICENSES = {LicenseType.MIT, LicenseType.APACHE_2, LicenseType.BSD}
573
574         def reclassify_licenses(self, args):
575             """Correct annotation-vs-video license conflation in existing rows.
576
577             A third-party video (one with a remote `video_url` ? i.e. broadcast footage like
578             ShuttleSet, whether or not it has since been downloaded to `local_path`) whose
579             `license_type` is a code/data license (MIT/Apache-2.0/BSD) is a conflation: that
580             license covers the ANNOTATIONS, not the footage. This moves the license to
581             `annotation_license`, sets the video `license_type` to Unknown, and recomputes
582             `is_commercial=False` so such videos stop being exported as commercially clean.
583
584             Dry-run by default; pass --apply to persist.
585             """
586             changed = []
587             seen = set()
588             for status in CurationStatus:
589                 for video in self.db.get_videos_by_status(status):
590                     if video.video_id in seen:
591                         continue
592                     seen.add(video.video_id)
593                     # The broadcast signal is a third-party video_url; a downloaded copy
594                     # (local_path set) does not change that the FOOTAGE rights are not ours.
595                     is_third_party = bool(video.video_url)
596                     if is_third_party and video.license_type in self._CODE_DATA_LICENSES:
597                         changed.append(video)
598
599             if not changed:
600                 print("No conflated rows found ? every remote video already has a video-
footage license.")
601                 return
602
603             print(f"'APPLYING' if args.apply else 'DRY-RUN' ? {len(changed)} conflated
row(s):")
604             for v in changed:
605                 print(f"    {v.video_id:40s} {v.license_type.value} (video) ->
annotation_license; "
606                     f"video license -> Unknown; is_commercial {v.is_commercial} -> False")
607                 if args.apply:
608                     v.annotation_license = v.annotation_license or v.license_type
609                     v.license_type = LicenseType.UNKNOWN
610                     v.is_commercial =
self.validator.check_license_commercial(v.license_type)
611                     self.db.insert_video(v)
612
613             if args.apply:
614                 print("\nApplied. Re-run `curate export` to regenerate the manifest without
these videos.")
615             else:
616                 print("\nDry-run only. Re-run with --apply to persist, then `curate

```



```

export`.")
617
618     def export_eval_set(self, args):
619         """Export curated annotations as indexer-ready eval/training sets.
620
621         Emits, per qualifying video, a `

```

```

678         # entry, so the downstream run knows the deterministic join isn't
available for it.
679         md5 = None
680         if video.local_path and os.path.exists(video.local_path):
681             md5 = md5_file(video.local_path)
682         elif video.local_path:
683             print(f"Warning: '{video.video_id}' local file missing
({video.local_path}); "
684                 f"exporting without an md5 (the indexer join needs the exact
file).")
685
686         manifest.append({
687             "video_id": video.video_id,
688             "sport": sport.value,
689             "csv": os.path.relpath(csv_path, out_dir).replace(os.sep, "/"),
690             "md5": md5,
691             "local_path": video.local_path,
692             "video_url": video.video_url,
693             "match_name": video.match_name,
694             "fps": video.fps,
695             "resolution": video.resolution,
696             "license_type": video.license_type.value,
697             "is_commercial": video.is_commercial,
698             "status": video.status.value,
699             "num_segments": len(intervals),
700         })
701
702     os.makedirs(out_dir, exist_ok=True)
703     manifest_doc = {
704         "commercial_only": commercial_only,
705         "statuses": [s.value for s in statuses],
706         "total_videos": len(manifest),
707         "total_segments": sum(m["num_segments"] for m in manifest),
708         "eval_sets": manifest,
709     }
710     # Validate the manifest against the typed contract before writing, so a
711     # field rename / type drift / miscount fails loudly at export time instead
712     # of silently breaking the downstream indexer's training/eval run.
713     EvalSetManifest.model_validate(manifest_doc).validate_counts()
714     manifest_path = os.path.join(out_dir, "manifest.json")
715     with open(manifest_path, "w", encoding="utf-8") as f:
716         json.dump(manifest_doc, f, indent=2)
717
718     print(f"\n=== Exported eval/training set to '{out_dir}' ===")
719     if manifest:
720         summary = [[m["sport"], m["video_id"], m["num_segments"],
721                     "YES" if m["is_commercial"] else "NO",
722                     os.path.basename(m["local_path"]) if m["local_path"] else "(no
local file)"]
723                   for m in manifest]
724         print(tabulate(summary, headers=["Sport", "Video ID", "Segments", "Comm?",
"Video file"], tablefmt="grid"))
725         print(f"\nVideos exported : {len(manifest)}")
726         print(f"Total segments : {manifest_doc['total_segments']}")
727         if skipped_noncomm:
728             print(f"Skipped (non-commercial) : {skipped_noncomm} (pass --include-
noncommercial to include)")
729         if skipped_empty:
730             print(f"Skipped (no annotations) : {skipped_empty}")
731         print(f"Manifest : {manifest_path}\n")
732
733     def prep_annotation_proxy(self, args):
734         """Re-encode a video into a frame-identical annotation proxy (ADR-002).
735
736         The proxy is what the annotation vendor receives: downscaled and

```

```

737 scrub-friendly (dense keyframes) so CVAT frame-step-1 work is cheap,
738 but with the source's EXACT frame count / fps / timeline, so frame N
739 on the proxy is frame N on the original and returned labels apply to
740 the original unchanged. Parity is verified after the encode; a proxy
741 that fails verification is deleted, never shared. Provenance (source +
742 proxy MD5s) is appended to the proxy manifest ? downstream processing
743 (import-local, export, the indexer) keeps running on the ORIGINAL file.
744 """
745 from src.core import proxy as proxylib
746
747 src = args.video_path
748 if not os.path.exists(src):
749     print(f"Error: video '{src}' not found.")
750     sys.exit(1)
751
752 stem = os.path.basename(src).split('.')[0]
753 # Mirror import-local's video_id derivation so the manifest row and a
754 # later DB registration of the same source agree on the key.
755 video_id = args.video_id or stem.lower().replace(" ", "_")
756 out = args.out or os.path.join("data", "videos", "proxies", f"{stem}_proxy.mp4")
757 # Same-file guard must survive case-insensitive filesystems (Windows):
758 # ffmpeg -y on a case-variant of the source would truncate the original.
759 same_file = (os.path.normcase(os.path.realpath(out))
760              == os.path.normcase(os.path.realpath(src)))
761 if not same_file and os.path.exists(out):
762     try:
763         same_file = os.path.samefile(out, src)
764     except OSError:
765         pass
766 if same_file:
767     print("Error: --out must differ from --video-path (it would overwrite "
768           "the original recording).")
769     sys.exit(1)
770 if os.path.exists(out) and not args.force:
771     print(f"Error: proxy '{out}' already exists; pass --force to overwrite.")
772     sys.exit(1)
773
774 encoder = getattr(args, "encoder", proxylib.DEFAULT_ENCODER)
775 if encoder not in proxylib.ENCODERS:
776     print(f"Error: unknown --encoder '{encoder}' (expected one of
777 {proxylib.ENCODERS}).")
778     sys.exit(1)
779 if not proxylib.encoder_available(encoder):
780     print(f"Error: encoder '{encoder}' is not available in this ffmpeg build "
781           f"(no h264_nvenc / no GPU?). Use --encoder libx264 (CPU).")
782     sys.exit(1)
783
784 try:
785     src_info = proxylib.probe_video(src)
786 except proxylib.ProxyError as e:
787     print(f"Error: {e}")
788     sys.exit(1)
789 if not src_info.avg_fps: # None or 0 (e.g. a cover-art/attached-pic stream)
790     print(f"Error: could not determine fps of '{src}' ? the frame<->time "
791           f"contract needs it.")
792     sys.exit(1)
793 if proxylib.is_vfr(src_info) and not args.allow_vfr:
794     print(f"Error: '{src}' looks variable-frame-rate "
795           f"(r={src_info.r_fps} vs avg={src_info.avg_fps}).")
796     print(" The vendor pipeline recomputes times as frame/fps, which is only "
797           "exact for constant-frame-rate sources. Re-run with --allow-vfr to "
798           "proceed anyway (timestamps stay identical to the original, but "
799           "frame/fps seconds are approximate for BOTH files).")
800     sys.exit(1)

```

```

801         # Hash the source up front: the provenance row needs it anyway, and it
802         # lets the manifest collision guard run BEFORE the expensive encode.
803         print("Hashing source (provenance)...")
804         source_md5 = proxylib.md5_file(src)
805         existing = proxylib.read_manifest_rows(args.manifest)
806         clashes = [r for r in existing
807                    if r.get("video_id") == video_id and r.get("source_md5") !=
source_md5]
808         if clashes and not args.force:
809             print(f"Error: manifest '{args.manifest}' already maps video_id "
810                   f"'{video_id}' to a DIFFERENT source (md5
{clashes[-1].get('source_md5')}).")
811             print("  Filename-derived video_ids collide easily ? pass a distinct "
812                   "--video-id to keep both, or --force to append anyway (last row
wins).")
813             sys.exit(1)
814         if any(r.get("video_id") == video_id and r.get("source_md5") == source_md5
815               for r in existing):
816             print(f"Note: manifest already has row(s) for '{video_id}' with this "
817                   f"source; the new row supersedes them (last row wins).")
818
819         os.makedirs(os.path.dirname(os.path.abspath(out)), exist_ok=True)
820         cmd = proxylib.build_proxy_command(src, out, height=args.height,
821                                           crf=args.crf, gop=args.gop,
822                                           preset=args.preset, encoder=encoder)
823         print(f"Encoding proxy ({src_info.resolution} -> height<={args.height}, "
824               f"f{fps} {src_info.avg_fps:g} preserved, encoder={encoder})...")
825         # From here on a failure must never leave a proxy file behind: a partial
826         # or unverified file under the to-be-shared name is exactly the artifact
827         # this flow exists to prevent (and -y may already have truncated a
828         # previous good proxy on --force).
829         try:
830             proxylib.run_ffmpeg(cmd)
831             proxy_info = proxylib.probe_video(out)
832         except proxylib.ProxyError as e:
833             print(f"Error: {e}")
834             self._delete_failed_proxy(out)
835             sys.exit(1)
836         except BaseException:
837             # KeyboardInterrupt mid-encode included ? clean up, then re-raise.
838             self._delete_failed_proxy(out)
839             raise
840
841         errors = proxylib.verify_parity(src_info, proxy_info)
842         if errors:
843             print("Error: proxy FAILED frame<->time parity verification:")
844             for err in errors:
845                 print(f"  - {err}")
846             self._delete_failed_proxy(out)
847             sys.exit(1)
848
849         print("Hashing proxy (provenance)...")
850         proxy_md5 = proxylib.md5_file(out)
851
852         from datetime import datetime, timezone
853         row = {
854             "video_id": video_id,
855             "source_path": os.path.abspath(src),
856             "source_md5": source_md5,
857             "proxy_path": os.path.abspath(out),
858             "proxy_md5": proxy_md5,
859             "fps": f"{src_info.avg_fps:g}",
860             "frame_count": src_info.nb_frames,
861             "duration_s": f"{src_info.duration:.3f}" if src_info.duration is not None
else "",

```

```

862         "source_resolution": src_info.resolution,
863         "proxy_resolution": proxy_info.resolution,
864         "encoder_settings": (
865             (f"h264_nvenc preset={proxylib.NVENC_PRESET} rc=vbr cq={args.crf}"
866              if encoder == "nvenc"
867              else f"libx264 crf={args.crf} preset={args.preset}")
868             + f" g={args.gop} height<={args.height} fps_mode=passthrough
audio=aac160k"),
869         "created_at": datetime.now(timezone.utc).strftime("%Y-%m-%dT%H:%M:%SZ"),
870     }
871     try:
872         proxylib.append_manifest_row(args.manifest, row)
873     except OSError as e:
874         # E.g. the CSV is open in Excel. The proxy itself is verified and
875         # fine ? surface the row so the provenance isn't lost, keep the file.
876         print(f"Error: could not append to manifest '{args.manifest}': {e}")
877         print(" The proxy IS verified and kept. Close the manifest file and "
878               "append this row manually (or re-run with --force):")
879         for k in proxylib.MANIFEST_FIELDS:
880             print(f"    {k}: {row.get(k, '')}")
881         sys.exit(1)
882
883     src_mb = os.path.getsize(src) / (1024 * 1024)
884     out_mb = os.path.getsize(out) / (1024 * 1024)
885     print("\n=== Annotation proxy ready ===")
886     print(f" Video ID      : {video_id}")
887     print(f" Source         : {src} ({src_info.resolution}, {src_mb:.0f} MB, md5
{source_md5})")
888     print(f" Proxy          : {out} ({proxy_info.resolution}, {out_mb:.0f} MB, md5
{proxy_md5})")
889     print(f" Parity         : VERIFIED ? {src_info.nb_frames} frames @
{src_info.avg_fps:g} fps on both")
890     print(f" Manifest        : {args.manifest}")
891     print("\nNext:")
892     print(" 1. Share the PROXY with the vendor (CVAT at frame step 1; first/last "
893           "box of every rally on the exact contact/dead frame).")
894     print(" 2. On delivery, run convert-cvat with --video-path pointing at the "
895           "ORIGINAL (or pass --fps from the manifest row).")
896     print(" 3. import-local / export keep pointing at the ORIGINAL file ? the "
897           "indexer keys videos by its MD5.\n")
898
899     def prep_annotation_batch(self, args):
900         """Mirror a tree of source videos into annotation proxies (ADR-002).
901
902         Input layout: ``<input-dir>/<match-folder>/<well-named video>`` (videos
903         sitting directly in ``<input-dir>`` are accepted too). The output
904         mirrors the SAME folder layout under ``--out-dir`` but holds ONLY the
905         proxies ? i.e. exactly the tree that gets shared with the vendor.
906
907         A **preflight gate** runs first (see ``src/core/batch_preflight.py``):
908         it probes/fingerprints every video and BLOCKS the batch on duplicates
909         (the same video in two folders) unless ``--allow-duplicates``, while
910         surfacing empty subdirs, VFR sources, and multi-video folders as
911         warnings. ``--dry-run`` prints that report and stops (no encode).
912
913         Every video then runs through the full single-video flow (probe, VFR
914         gate, parity verification, provenance manifest, all guards); one
915         video's failure is reported and does not stop the rest. Already-proxied
916         videos are skipped without --force, so re-runs are cheap and resumable.
917         """
918         from src.core import batch_preflight as pf
919
920         root = args.input_dir
921         if not os.path.isdir(root):
922             print(f"Error: --input-dir '{root}' is not a directory.")

```

```

923         sys.exit(1)
924
925     # --- Preflight gate (probe + fingerprint the whole tree, once) -----
926     report = pf.run_preflight(root)
927     print(pf.render_preflight(report))
928
929     if report.n_videos == 0:
930         print(f"\nError: no video files found under '{root}' "
931               f"(recognized: {'', '.join(sorted(pf.VIDEO_EXTS))}).")
932         sys.exit(1)
933
934     if report.has_blockers and not getattr(args, "allow_duplicates", False):
935         print("\nGATE: BLOCKED ? duplicate source videos detected (see above). "
936               "Remove the copy, or pass --allow-duplicates if intentional.")
937         sys.exit(1)
938     if report.has_blockers:
939         print("\nNote: --allow-duplicates given ? proceeding despite duplicate
sources.")
940
941     if getattr(args, "dry_run", False):
942         print("\n--dry-run: preflight only, no proxies written.")
943         return
944
945     # Encode targets come straight from the preflight (already probed).
946     targets = [(e.rel_dir, e.path) for e in report.entries]
947     results = []
948     for idx, (rel, src) in enumerate(targets, start=1):
949         stem = os.path.basename(src).split('.')[0]
950         out = os.path.join(args.out_dir, rel, f"{stem}_proxy.mp4")
951         label = os.path.join(rel, os.path.basename(src)) if rel else
os.path.basename(src)
952         if os.path.exists(out) and not args.force:
953             results.append((rel or ".", os.path.basename(src), "have",
954                             "proxy exists (--force to redo)"))
955             continue
956         one = argparse.Namespace(
957             video_path=src, out=out, video_id=None,
958             height=args.height, crf=args.crf, gop=args.gop, preset=args.preset,
959             encoder=getattr(args, "encoder", "libx264"),
960             manifest=args.manifest, allow_vfr=args.allow_vfr, force=args.force,
961         )
962         print(f"\n--- [{idx}/{len(targets)}] {label} ---")
963         try:
964             self.prep_annotation_proxy(one)
965             results.append((rel or ".", os.path.basename(src), "ok",
966                             os.path.relpath(out, args.out_dir)))
967         except SystemExit:
968             # The single-video flow already printed the error and cleaned up.
969             results.append((rel or ".", os.path.basename(src), "FAIL",
970                             "see error above"))
971
972     print("\n=== Batch proxy summary ===")
973     print(tabulate([[r, v, s, d] for r, v, s, d in results],
974                   headers=["Folder", "Video", "Status", "Detail"],
975                   tablefmt="grid"))
976     ok = sum(1 for _, _, s, _ in results if s == "ok")
977     have = sum(1 for _, _, s, _ in results if s == "have")
978     failed = [r for r in results if r[2] == "FAIL"]
979     print(f"\nProxied : {ok}   Already-have : {have}   Failed : {len(failed)}")
980     print(f"Manifest: {args.manifest}")
981     print(f"Share the '{args.out_dir}' tree with the vendor ? it contains ONLY "
982           f"proxies; every other stage runs on the originals.\n")
983     if failed:
984         sys.exit(1)

```

```

985     @staticmethod
986     def _delete_failed_proxy(path: str):
987         """Remove a failed/unverified proxy ? it must never reach the vendor."""
988         if not os.path.exists(path):
989             return
990         try:
991             os.remove(path)
992             print(f"Deleted the failed proxy '{path}' (an unverified proxy must "
993                   f"never be shared ? it poisons golden labels).")
994         except OSError:
995             print(f"WARNING: could not delete the failed proxy '{path}' ? do NOT "
996                   f"share it with the vendor.")
997
998     def convert_cvat(self, args):
999         """Convert a vendor CVAT rally export into our canonical per-rally CSV.
1000
1001         Absorbs the CVAT bounding-box delivery format (rally fields carried as box
1002         attributes) on the ingestion side: each rally's time is recomputed from
1003         its frame range at the video's true fps, and content defects (overlaps,
1004         zero-duration, off-vocabulary endings, coarse sampling) are written to an
1005         issues report rather than silently "fixed". The emitted CSV imports via
1006         `import-local`; any blocker issue must be resolved first, since the
1007         importer's validator rejects overlapping / zero-duration rallies.
1008         """
1009         from src.parsers.cvat.rally_cvat import (
1010             convert, write_canonical_csv, render_issue_report,
1011         )
1012         try:
1013             result = convert(args.input, fps=args.fps,
1014                             video_path=args.video_path, step=args.step)
1015         except (FileNotFoundError, ValueError) as e:
1016             print(f"Error: {e}")
1017             sys.exit(1)
1018
1019         if not result.records:
1020             print(f"Error: the CVAT export parsed to ZERO rallies ({args.input}). "
1021                   f"Likely the wrong job was exported or the rally_number attribute was "
1022                   f"renamed (boxes without a matched rally_number are skipped).")
1023             sys.exit(1)
1024
1025         write_canonical_csv(result.records, args.out)
1026
1027         # Observability report next to the canonical CSV. Never raises.
1028         from src.reporting import emit_delivery_report
1029         src_name = os.path.splitext(os.path.basename(args.input))[0]
1030         dur = (result.meta.get("stop_frame") / result.fps) if
1031         (result.meta.get("stop_frame") and result.fps) else None
1032         emit_delivery_report(
1033             out_dir=os.path.dirname(os.path.abspath(args.out)) or ".", source=src_name,
1034             records=result.records, issues=result.issues, fps=result.fps,
1035             video_path=args.video_path,
1036             video_duration_s=dur, source_format="cvat", step=result.step,
1037         )
1038
1039         report = render_issue_report(result, os.path.basename(args.input))
1040         if args.issues:
1041             with open(args.issues, "w", encoding="utf-8") as f:
1042                 f.write(report)
1043         print(report)
1044         print(f"Canonical CSV written: {args.out} ({len(result.records)} rallies)")
1045         if result.n_blockers:
1046             print(f"\n {result.n_blockers} BLOCKER issue(s) ? import-local will reject "
1047                   f"this file until they are resolved (by the vendor or manual

```

```

review).")
1046         else:
1047             print("\n No blockers. Next:")
1048             print(f"    python -m src.cli.curate import-local --sport <sport> "
1049                   f"--video-path <video> --annotations-path {args.out} --fps
{result.fps:g}")
1050
1051     def validate_delivery(self, args):
1052         """Run the full validation suite on a rally delivery and emit a report.
1053
1054         Accepts a CVAT export (--input + --video-path/--fps) or an already-canonical
1055         CSV (--csv). Writes, into --out-dir:
1056             <id>_report.md    human report (scoreboard + issues + rallies to review)
1057             <id>_report.json  machine summary (for CI / a bulk-ingest loop)
1058             <id>_review.csv   the annotatable manual-verification sheet
1059         Exits non-zero if any blocker issue exists, so it can gate a bulk pipeline.
1060         """
1061         import json
1062         from src.parsers.cvat.rally_cvat import (
1063             convert, load_canonical_csv, detect_issues,
1064             build_summary, render_validation_md, write_review_csv,
1065         )
1066         vid = args.video_id or "delivery"
1067         fps = None
1068         step = None
1069         stop_frame = None
1070         if args.input:
1071             try:
1072                 result = convert(args.input, fps=args.fps, video_path=args.video_path)
1073             except (FileNotFoundError, ValueError) as e:
1074                 print(f"Error: {e}")
1075                 sys.exit(1)
1076             records, issues, fps = result.records, result.issues, result.fps
1077             step = result.step
1078             stop_frame = result.meta.get("stop_frame")
1079             if not records:
1080                 print(f"Error: the CVAT export parsed to ZERO rallies ({args.input}). "
1081                       f"The wrong job was exported or the rally_number attribute was
renamed. "
1082                       f"GATE: BLOCKED.")
1083                 sys.exit(2)
1084         elif args.csv:
1085             try:
1086                 records = load_canonical_csv(args.csv)
1087             except FileNotFoundError as e:
1088                 print(f"Error: {e}")
1089                 sys.exit(1)
1090             if not records:
1091                 print(f"Error: no usable rally rows in {args.csv}")
1092                 sys.exit(1)
1093             issues = detect_issues(records)
1094         else:
1095             print("Error: provide --input <cvat.xml> or --csv <canonical.csv>")
1096             sys.exit(1)
1097
1098         os.makedirs(args.out_dir, exist_ok=True)
1099         summary = build_summary(records, issues, source=vid, fps=fps)
1100         md = render_validation_md(records, issues, source=vid)
1101         json_path = os.path.join(args.out_dir, f"{vid}_report.json")
1102         md_path = os.path.join(args.out_dir, f"{vid}_report.md")
1103         review_path = os.path.join(args.out_dir, f"{vid}_review.csv")
1104         with open(json_path, "w", encoding="utf-8") as f:
1105             json.dump(summary, f, indent=2)
1106         with open(md_path, "w", encoding="utf-8") as f:
1107             f.write(md)

```



```

1108         write_review_csv(records, issues, review_path)
1109
1110         # Observability report (envelope JSON + technical md + customer summary),
1111         # colocated with the existing artifacts. Never raises.
1112         from src.reporting import emit_delivery_report
1113         duration_s = (stop_frame / fps) if (stop_frame and fps) else None
1114         report_paths = emit_delivery_report(
1115             out_dir=args.out_dir, source=vid, records=records, issues=issues, fps=fps,
1116             video_path=args.video_path, video_duration_s=duration_s,
1117             source_format=("cvat" if args.input else "csv"), step=step,
1118         )
1119
1120         print(md)
1121         print(f"\nWrote:\n {md_path}\n {json_path}\n {review_path}")
1122         if report_paths:
1123             print(f" {report_paths.get('technical_md')}\n
1124 {report_paths.get('customer_md')}\n {report_paths.get('json')}")
1125         if summary["n_blockers"]:
1126             print(f"\n GATE: BLOCKED ? {summary['n_blockers']} blocker(s) must be
1127 resolved.")
1128             sys.exit(2)
1129         print("\n GATE: PASS ? review the flagged rallies, then `import-local`.")
1130
1131     def curate(self):
1132         """Interactive loop to review and curate staging video annotations."""
1133         staging_videos = self.db.get_videos_by_status(CurationStatus.STAGING)
1134         if not staging_videos:
1135             print("\nNo staging video annotations to review.\n")
1136             return
1137
1138         print(f"\nFound {len(staging_videos)} matches in STAGING state for curation:")
1139         table_data = []
1140         for idx, video in enumerate(staging_videos):
1141             table_data.append([
1142                 idx + 1,
1143                 video.video_id,
1144                 video.match_name,
1145                 video.sport.value,
1146                 video.license_type.value,
1147                 "YES" if video.is_commercial else "NO"
1148             ])
1149
1150         print(tabulate(table_data, headers=["#", "Video ID", "Match Name", "Sport",
1151 "License", "Commercial Safe?"], tablefmt="grid"))
1152         print("\nEnter a number (1-{}) to review details, or 'q' to
1153 quit:".format(len(staging_videos)))
1154
1155         choice = input("> ").strip()
1156         if choice.lower() == 'q':
1157             return
1158
1159         try:
1160             idx = int(choice) - 1
1161             if idx < 0 or idx >= len(staging_videos):
1162                 print("Invalid choice.")
1163                 return
1164             self._curate_video_interactive(staging_videos[idx])
1165         except ValueError:
1166             print("Invalid input.")
1167
1168     def _curate_video_interactive(self, video: VideoMetadata):
1169         """Displays details of a specific staging video and prompts for curation
1170 action."""
1171
1172         # Load and count events

```

24. user (2026-07-02T04:49:01.196Z)

```
1168         rallies = []
1169         pts = []
1170         delivs = []
1171         if video.sport == SportType.BADMINTON:
1172             rallies = self.db.get_badminton_rallies(video.video_id)
1173         elif video.sport == SportType.PICKLEBALL:
1174             rallies = self.db.get_pickleball_rallies(video.video_id)
1175         elif video.sport == SportType.TABLE_TENNIS:
1176             rallies = self.db.get_tabletennis_rallies(video.video_id)
1177         elif video.sport == SportType.TENNIS:
1178             pts = self.db.get_tennis_points(video.video_id)
1179         elif video.sport == SportType.CRICKET:
1180             delivs = self.db.get_cricket_deliveries(video.video_id)
1181
1182         while True:
1183             print("\n===== REVIEWING VIDEO =====")
1184             print(f"Match Name      : {video.match_name}")
1185             print(f"Video ID       : {video.video_id}")
1186             print(f"Sport          : {video.sport.value}")
1187             print(f"Video URL      : {video.video_url or 'N/A'}")
1188             print(f"Local Path     : {video.local_path or 'N/A'}")
1189             print(f"License       : {video.license_type.value} ({video.license_url or 'no
URL'})")
1190             print(f"Commercial OK: {video.is_commercial}")
1191
1192             if video.sport in (SportType.BADMINTON, SportType.PICKLEBALL,
SportType.TABLE_TENNIS):
1193                 print(f"Total Rallies: {len(rallies)}")
1194                 if rallies:
1195                     print("Sample rallies:")
1196                     sample_data = [[r.rally_number, f"{r.start_time}s - {r.end_time}s",
r.score_after, r.shots_count] for r in rallies[:5]]
1197                     print(tabulate(sample_data, headers=["Rally #", "Duration", "Score",
"Shots"], tablefmt="simple"))
1198                 elif video.sport == SportType.TENNIS:
1199                     print(f"Total Points : {len(pts)}")
1200                     if pts:
1201                         print("Sample points:")
1202                         sample_data = [[p.point_number, f"{p.start_time}s - {p.end_time}s",
p.game_score] for p in pts[:5]]
1203                         print(tabulate(sample_data, headers=["Point #", "Duration",
"Score"], tablefmt="simple"))
1204                 elif video.sport == SportType.CRICKET:
1205                     print(f"Total Balls  : {len(delivs)}")
1206                     if delivs:
1207                         print("Sample deliveries:")
1208                         sample_data = [[f"{d.over}.{d.ball}", f"{d.start_time}s -
{d.end_time}s", d.runs, "Wicket" if d.wicket else ""] for d in delivs[:5]]
1209                         print(tabulate(sample_data, headers=["Over.Ball", "Duration",
"Runs", "Wicket"], tablefmt="simple"))
1210
1211                 print("\nChoose Curation Action:")
1212                 print(" [a] Approve (promote to CURATED)")
1213                 print(" [r] Reject (delete from database)")
1214                 print(" [p] Play a rally/point segment in browser")
1215                 print(" [k] Keep in Staging (do nothing and go back)")
1216
1217                 action = input("> ").strip().lower()
1218                 if action == 'a':
1219                     self.db.update_status(video.video_id, CurationStatus.CURATED)
1220                     print(f"Successfully promoted '{video.match_name}' to CURATED.\n")
1221                     break
1222                 elif action == 'r':
```

```

1223         self.db.delete_video(video.video_id)
1224         print(f"Successfully rejected and deleted '{video.match_name}'.\n")
1225         break
1226     elif action == 'k':
1227         print("Curation state unchanged.\n")
1228         break
1229     elif action == 'p':
1230         self._play_segment_interactive(video, rallies, pts, delivs)
1231     else:
1232         print("Invalid action selected.")
1233
1234     def _play_segment_interactive(self, video: VideoMetadata, rallies:
List[BadmintonRally], pts: List[TennisPoint], delivs: List[CricketDelivery]):
1235         import webbrowser
1236
1237         # Get start/end based on user input
1238         start_time = None
1239         end_time = None
1240         label = "segment"
1241
1242         if video.sport in (SportType.BADMINTON, SportType.PICKLEBALL,
SportType.TABLE_TENNIS):
1243             if not rallies:
1244                 print("No rallies to play.")
1245                 return
1246             idx_str = input(f"Enter Rally number to play (1-{len(rallies)}): ").strip()
1247             try:
1248                 rally_num = int(idx_str)
1249                 rally = next((r for r in rallies if r.rally_number == rally_num), None)
1250                 if rally:
1251                     start_time = rally.start_time
1252                     end_time = rally.end_time
1253                     label = f"Rally {rally_num}"
1254             except ValueError:
1255                 pass
1256         elif video.sport == SportType.TENNIS:
1257             if not pts:
1258                 print("No points to play.")
1259                 return
1260             idx_str = input(f"Enter Point number to play (1-{len(pts)}): ").strip()
1261             try:
1262                 pt_num = int(idx_str)
1263                 pt = next((p for p in pts if p.point_number == pt_num), None)
1264                 if pt:
1265                     start_time = pt.start_time
1266                     end_time = pt.end_time
1267                     label = f"Point {pt_num}"
1268             except ValueError:
1269                 pass
1270         elif video.sport == SportType.CRICKET:
1271             if not delivs:
1272                 print("No deliveries to play.")
1273                 return
1274             over_str = input("Enter Over number: ").strip()
1275             ball_str = input("Enter Ball number: ").strip()
1276             try:
1277                 o_num = int(over_str)
1278                 b_num = int(ball_str)
1279                 d = next((x for x in delivs if x.over == o_num and x.ball == b_num),
None)
1280                 if d:
1281                     start_time = d.start_time
1282                     end_time = d.end_time
1283                     label = f"Over {o_num} Ball {b_num}"
1284             except ValueError:

```

```

1285         pass
1286
1287     if start_time is None or end_time is None:
1288         print("Segment not found or invalid input.")
1289         return
1290
1291     print(f"Playing {label} (start: {start_time}s, end: {end_time}s)...")
1292
1293     # Check video location
1294     play_url = None
1295     if video.local_path and os.path.exists(video.local_path):
1296         abs_path = os.path.abspath(video.local_path)
1297         # Create file:// URL with media fragment (browsers support #t=start,end
fragment)
1298         play_url = f"file:/// {abs_path.replace(os.sep,
1299         '/' )}#t={start_time},{end_time}"
1300     elif video.video_url:
1301         # YouTube URL: seek to start time
1302         if "youtube.com" in video.video_url or "youtu.be" in video.video_url:
1303             t_sec = int(start_time)
1304             sep = "&" if "?" in video.video_url else "?"
1305             play_url = f"{video.video_url}{sep}t={t_sec}s"
1306         else:
1307             play_url = video.video_url
1308
1309     if play_url:
1310         print(f"Opening browser: {play_url}")
1311         webbrowser.open(play_url)
1312     else:
1313         print("Error: No local video file or remote URL available to play.")
1314
1315 def main():
1316     parser = argparse.ArgumentParser(description="Sports Data Curator CLI")
1317     subparsers = parser.add_subparsers(dest="command")
1318
1319     # Status command
1320     subparsers.add_parser("status", help="Display summary status of sports metadata
database")
1321
1322     # Curate command
1323     subparsers.add_parser("curate", help="Launch interactive curation tool for staging
entries")
1324
1325     # Import-local command
1326     import_parser = subparsers.add_parser("import-local", help="Manually import a local
video and annotations")
1327     import_parser.add_argument("--sport", required=True, help="Sport type (e.g.
badminton, tennis, cricket)")
1328     import_parser.add_argument("--video-path", required=True, help="Path to local video
file")
1329     import_parser.add_argument("--annotations-path", required=True, help="Path to local
JSON/CSV annotations file")
1330     import_parser.add_argument("--match-name", help="Name of the match (defaults to
video filename)")
1331     import_parser.add_argument("--video-id", help="Custom alphanumeric video ID")
1332     import_parser.add_argument("--license", default="Proprietary", help="License type
(e.g. MIT, CC-BY-NC, Proprietary)")
1333     import_parser.add_argument("--fps", type=float, default=30.0, help="Frames per
second of the video")
1334     import_parser.add_argument("--resolution", default="1080p", help="Resolution of the
video (e.g., 1080p, 720p)")
1335     import_parser.add_argument("--force", action="store_true",
1336         help="Overwrite an existing video's annotations (default:
refuse, to protect golden labels)")

```

```

1337     import_parser.add_argument("--normalize", action="store_true",
1338                               help="Run normalize-annotations on a CSV input first
(drops unusable rows, "
1339                               "renumbers fractional/duplicate rally ids 1..N by
start time)")
1340
1341     # Fetch command ? (re)download videos at best resolution, update DB in place
1342     fetch_parser = subparsers.add_parser(
1343         "fetch", help="(Re)download videos at best available resolution, updating the DB
in place")
1344     fetch_parser.add_argument("--sport", help="Restrict to a single sport (default: all
supported_sports)")
1345     fetch_parser.add_argument("--video-id", help="Fetch a single video by ID")
1346     fetch_parser.add_argument("--status", choices=["staging", "curated", "all"],
default="all",
1347                               help="Which curation statuses to fetch (default: all)")
1348     fetch_parser.add_argument("--min-height", type=int, default=720,
1349                               help="Skip videos already at/above this height unless
--force (default: 720)")
1350     fetch_parser.add_argument("--max-height", type=int, default=None,
1351                               help="Cap download resolution (e.g. 1080); default: best
available")
1352     fetch_parser.add_argument("--force", action="store_true",
1353                               help="Re-download even if a good-enough local file already
exists")
1354     fetch_parser.add_argument("--include-noncommercial", action="store_true",
1355                               help="Include videos whose license is not commercially
safe (default: commercial only)")
1356     fetch_parser.add_argument("--cookies-from-browser", dest="cookies_from_browser",
metavar="BROWSER",
1357                               help="Use a logged-in browser's YouTube cookies (e.g.
chrome, edge, firefox, brave). "
1358                               "A Premium login gives reliable best-resolution
downloads. "
1359                               "Chromium browsers must be CLOSED on Windows (cookie
DB lock).")
1360     fetch_parser.add_argument("--cookies", metavar="FILE",
1361                               help="Path to a Netscape-format cookies.txt (alternative
to --cookies-from-browser)")
1362     fetch_parser.add_argument("--player-client", dest="player_client",
default="default", metavar="CLIENT",
1363                               help="YouTube player client yt-dlp uses (default:
'default'). Avoid web/web_safari "
1364                               "(SABR-forced -> drops to 144p). 'tv'/'mweb' also
serve 1080p if needed.")
1365     fetch_parser.add_argument("--sleep", type=float, default=4.0, metavar="SECONDS",
1366                               help="Pause between videos to avoid YouTube rate-limiting
(default: 4)")
1367     fetch_parser.add_argument("--retries", type=int, default=2, metavar="N",
1368                               help="Retries per video when a download fails or returns a
low format (default: 2)")
1369     fetch_parser.add_argument("--retry-backoff", type=float, default=20.0,
metavar="SECONDS",
1370                               help="Base backoff between retries; grows linearly per
attempt (default: 20)")
1371
1372     # Prep-annotation-proxy command ? vendor-handoff re-encode (ADR-002)
1373     prep_parser = subparsers.add_parser(
1374         "prep-annotation-proxy",
1375         help="Re-encode a video into a light, frame-identical annotation proxy for "
1376         "vendor handoff (ADR-002): downscaled + scrub-friendly, same frame "
1377         "count/fps/timeline, parity-verified, provenance recorded")
1378     prep_parser.add_argument("--video-path", required=True,
1379                               help="The ORIGINAL source video (stays the system-of-record
file)")

```

```

1380     prep_parser.add_argument("--out",
1381                               help="Proxy output path (default:
data/videos/proxies/<stem>_proxy.mp4)")
1382     prep_parser.add_argument("--video-id",
1383                               help="Manifest key (default: derived from the filename,
like import-local)")
1384     prep_parser.add_argument("--height", type=int, default=1080,
1385                               help="Max proxy height; never upscales (default: 1080 ? "
1386                                     "720 risks losing the table-tennis ball)")
1387     prep_parser.add_argument("--crf", type=int, default=19,
1388                               help="x264 quality, lower=better (default: 19)")
1389     prep_parser.add_argument("--gop", type=int, default=15,
1390                               help="Max keyframe interval in frames; small = snappy "
1391                                     "frame-stepping (default: 15)")
1392     prep_parser.add_argument("--preset", default="medium",
1393                               help="libx264 speed/size preset (default: medium; ignored
for nvenc)")
1394     prep_parser.add_argument("--encoder", choices=["libx264", "nvenc"],
default="libx264",
1395                               help="Video encoder: libx264 (CPU, default) or nvenc "
1396                                     "(NVIDIA GPU ? much faster on 4K, slightly larger
files)")
1397     prep_parser.add_argument("--manifest", default=os.path.join("data",
"annotation_proxy_manifest.csv"),
1398                               help="Provenance CSV to append to (default:
data/annotation_proxy_manifest.csv; "
1399                                     "commit it ? it is the proxy<->original mapping of
record)")
1400     prep_parser.add_argument("--allow-vfr", action="store_true",
1401                               help="Proceed on a variable-frame-rate source (frame/fps "
1402                                     "seconds are approximate; default: refuse)")
1403     prep_parser.add_argument("--force", action="store_true",
1404                               help="Overwrite an existing proxy file / append a manifest
"
1405                                     "row for a video_id already mapped to a different
source")
1406
1407     # Prep-annotation-batch command ? proxy a whole directory tree (ADR-002)
1408     batch_parser = subparsers.add_parser(
1409         "prep-annotation-batch",
1410         help="Proxy every video under a directory tree, mirroring its folder layout "
1411               "into --out-dir (the proxies-only tree that goes to the vendor)")
1412     batch_parser.add_argument("--input-dir", required=True,
1413                               help="Root directory: each subdirectory holds well-named "
1414                                     "source video(s); root-level videos also accepted")
1415     batch_parser.add_argument("--out-dir", required=True,
1416                               help="Mirrored output root ? receives ONLY the proxies")
1417     batch_parser.add_argument("--height", type=int, default=1080,
1418                               help="Max proxy height; never upscales (default: 1080)")
1419     batch_parser.add_argument("--crf", type=int, default=19,
1420                               help="x264 quality, lower=better (default: 19)")
1421     batch_parser.add_argument("--gop", type=int, default=15,
1422                               help="Max keyframe interval in frames (default: 15)")
1423     batch_parser.add_argument("--preset", default="medium",
1424                               help="libx264 speed/size preset (default: medium; ignored
for nvenc)")
1425     batch_parser.add_argument("--encoder", choices=["libx264", "nvenc"],
default="libx264",
1426                               help="Video encoder: libx264 (CPU, default) or nvenc "
1427                                     "(NVIDIA GPU ? much faster on 4K, slightly larger
files)")
1428     batch_parser.add_argument("--manifest", default=os.path.join("data",
"annotation_proxy_manifest.csv"),
1429                               help="Provenance CSV appended per video (default: "
1430                                     "data/annotation_proxy_manifest.csv)")

```

```

1431     batch_parser.add_argument("--allow-vfr", action="store_true",
1432                               help="Proceed on variable-frame-rate sources (default:
refuse)")
1433     batch_parser.add_argument("--allow-duplicates", action="store_true",
1434                               help="Proceed even if the same video appears in >1 folder
"
1435                               "(default: BLOCK; use for an intentional A/A test)")
1436     batch_parser.add_argument("--dry-run", action="store_true",
1437                               help="Run the preflight scan and print the report, then
stop "
1438                               "(no proxies written)")
1439     batch_parser.add_argument("--force", action="store_true",
1440                               help="Redo existing proxies / re-key colliding manifest
video_ids")
1441
1442     # Convert-cvat command ? normalize a vendor CVAT export into our canonical CSV
1443     cvat_parser = subparsers.add_parser(
1444         "convert-cvat",
1445         help="Convert a vendor CVAT rally export (bounding-box XML) into the canonical
per-rally CSV")
1446     cvat_parser.add_argument("--input", required=True,
1447                              help="Path to the CVAT XML export (image- or track-mode)")
1448     cvat_parser.add_argument("--out", required=True,
1449                              help="Path to write the canonical per-rally CSV (import-
local input)")
1450     cvat_parser.add_argument("--video-path",
1451                              help="Source video; its fps is read via ffprobe (seconds =
frame / fps)")
1452     cvat_parser.add_argument("--fps", type=float,
1453                              help="Override fps instead of probing the video")
1454     cvat_parser.add_argument("--step", type=int,
1455                              help="Frame sampling step (default: read from CVAT meta,
else inferred)")
1456     cvat_parser.add_argument("--issues",
1457                              help="Path to also write the human-readable issues report")
1458
1459     # Validate-delivery command ? run the full check suite + emit a structured report
1460     vd_parser = subparsers.add_parser(
1461         "validate-delivery",
1462         help="Validate a rally delivery (CVAT export or canonical CSV) and emit report +
review sheet")
1463     vd_parser.add_argument("--input", help="CVAT XML export (use with --video-
path/--fps)")
1464     vd_parser.add_argument("--csv", help="An already-canonical per-rally CSV to
validate")
1465     vd_parser.add_argument("--video-path",
1466                              help="Source video; fps via ffprobe (when using --input)")
1467     vd_parser.add_argument("--fps", type=float, help="Override fps (when using
--input)")
1468     vd_parser.add_argument("--video-id", help="Names the report files (<id>_report.md,
etc.)")
1469     vd_parser.add_argument("--out-dir", default="data/validation",
1470                              help="Where to write the report + review sheet")
1471
1472     # Export command
1473     export_parser = subparsers.add_parser(
1474         "export", help="Export curated annotations as indexer-ready eval/training CSVs +
manifest")
1475     export_parser.add_argument("--out-dir", default="data/eval_export",
1476                              help="Output directory for per-video CSVs and
manifest.json")
1477     export_parser.add_argument("--sport",
1478                              help="Restrict to a single sport (default: all
supported_sports)")
1479     export_parser.add_argument("--include-staging", action="store_true",

```

```

1480                 help="Also export videos still in staging (default:
exported only)")
1481     export_parser.add_argument("--include-noncommercial", action="store_true",
1482                               help="Include videos whose license is not commercially
safe (default: commercial only)")
1483
1484     # Reclassify-licenses command (one-shot corrective maintenance)
1485     reclass_parser = subparsers.add_parser(
1486         "reclassify-licenses",
1487         help="Fix annotation-vs-video license conflation (e.g. MIT stamped on broadcast
video)")
1488     reclass_parser.add_argument("--apply", action="store_true",
1489                                help="Persist changes (default: dry-run preview only)")
1490
1491     args = parser.parse_args()
1492     cli = SportsDataCLI()
1493
1494     try:
1495         _dispatch(cli, args, parser)
1496     except CurationError as e:
1497         # A library-style pre-flight failure (A2): print the actionable message and
1498         # exit 1, preserving the prior CLI behaviour now that _guard_import raises.
1499         print(f"Error: {e}")
1500         sys.exit(1)
1501
1502
1503     def _dispatch(cli: "SportsDataCLI", args, parser) -> None:
1504         """Route the parsed command to its handler. Separated so the CurationError
1505         catch in main() wraps every command uniformly."""
1506         if args.command == "status":
1507             cli.status()
1508         elif args.command == "reclassify-licenses":
1509             cli.reclassify_licenses(args)
1510         elif args.command == "curate":
1511             cli.curate()
1512         elif args.command == "import-local":
1513             cli.import_local(args)
1514         elif args.command == "prep-annotation-proxy":
1515             cli.prep_annotation_proxy(args)
1516         elif args.command == "prep-annotation-batch":
1517             cli.prep_annotation_batch(args)
1518         elif args.command == "convert-cvat":
1519             cli.convert_cvat(args)
1520         elif args.command == "validate-delivery":
1521             cli.validate_delivery(args)
1522         elif args.command == "export":
1523             cli.export_eval_set(args)
1524         elif args.command == "fetch":
1525             cli.fetch(args)
1526         else:
1527             # Default: print help
1528             parser.print_help()
1529
1530     if __name__ == "__main__":
1531         main()
1532

```

25. assistant (2026-07-02T04:49:07.771Z)

Now let me read the DB layer and schemas.

26. user (2026-07-02T04:49:10.195Z)

```

1     import sqlite3
2     import os

```



```

3     import logging
4     from contextlib import contextmanager
5     from enum import Enum
6     from typing import List, Dict, Any, Optional, Type, TypeVar
7     from src.core.schemas.base import VideoMetadata, SportType, CurationStatus, LicenseType
8     from src.core.schemas.badminton import BadmintonRally
9     from src.core.schemas.tennis import TennisPoint
10    from src.core.schemas.cricket import CricketDelivery
11    from src.core.schemas.pickleball import PickleballRally
12    from src.core.schemas.tabletennis import TableTennisRally
13
14    logger = logging.getLogger(__name__)
15
16    # Maps each sport to its annotation table and the column(s) that order segments
17    # chronologically. All child tables expose (start_time, end_time), so a
18    # single generic reader can serve every sport ? adding a new sport is one entry
19    # here plus its CREATE TABLE, with no new query code.
20    SEGMENT_TABLES = {
21        SportType.BADMINTON: ("badminton_rallies", "rally_number"),
22        SportType.TENNIS: ("tennis_points", "point_number"),
23        SportType.CRICKET: ("cricket_deliveries", "over_num, ball_num"),
24        SportType.PICKLEBALL: ("pickleball_rallies", "rally_number"),
25        SportType.TABLE_TENNIS: ("tabletennis_rallies", "rally_number"),
26    }
27
28
29    class SportsDatabase:
30        def __init__(self, db_path: str):
31            self.db_path = db_path
32            # Ensure target directory exists
33            os.makedirs(os.path.dirname(os.path.abspath(db_path)), exist_ok=True)
34            self._init_db()
35
36        @contextmanager
37        def _get_connection(self):
38            """Yield a SQLite connection and always close it.
39
40            Used as ``with self._get_connection() as conn:``. A plain
41            ``sqlite3.connect(...)`` used as a context manager commits/rolls back the
42            transaction but does NOT close the connection, which leaks handles and
43            raises ResourceWarnings. This wrapper guarantees the connection is closed
44            on exit; writers still commit explicitly inside the block.
45            """
46            conn = sqlite3.connect(self.db_path)
47            conn.row_factory = sqlite3.Row
48            # SQLite disables foreign keys per-connection by default; enable so the
49            # ON DELETE CASCADE declarations on the child tables actually fire.
50            conn.execute("PRAGMA foreign_keys = ON")
51            # Wait (up to 5s) for a competing writer's lock instead of failing immediately
52            with
53                # "database is locked" ? makes concurrent access robust (a sibling CLI, or a
54                # web-driven import_annotations writer) rather than racy. Writers still commit
55                # explicitly.
56                conn.execute("PRAGMA busy_timeout = 5000")
57                try:
58                    yield conn
59                finally:
60                    conn.close()
61
62        def _init_db(self):
63            with self._get_connection() as conn:
64                cursor = conn.cursor()
65
66                # 1. Create common videos table

```

```

65 cursor.execute("""
66 CREATE TABLE IF NOT EXISTS videos (
67     video_id TEXT PRIMARY KEY,
68     sport TEXT NOT NULL,
69     video_url TEXT,
70     local_path TEXT,
71     license_type TEXT NOT NULL,
72     license_url TEXT,
73     annotation_license TEXT,
74     is_commercial INTEGER NOT NULL DEFAULT 0,
75     status TEXT NOT NULL DEFAULT 'staging',
76     match_name TEXT NOT NULL,
77     fps REAL,
78     resolution TEXT
79 )
80 """)
81
82 # 2. Create badminton_rallies table
83 cursor.execute("""
84 CREATE TABLE IF NOT EXISTS badminton_rallies (
85     video_id TEXT NOT NULL,
86     rally_number INTEGER NOT NULL,
87     start_time REAL NOT NULL,
88     end_time REAL NOT NULL,
89     score_before TEXT,
90     score_after TEXT,
91     shots_count INTEGER,
92     ending_reason TEXT,
93     last_shot_outcome TEXT,
94     PRIMARY KEY (video_id, rally_number),
95     FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE
96 )
97 """)
98
99 # 3. Create tennis_points table
100 cursor.execute("""
101 CREATE TABLE IF NOT EXISTS tennis_points (
102     video_id TEXT NOT NULL,
103     point_number INTEGER NOT NULL,
104     start_time REAL NOT NULL,
105     end_time REAL NOT NULL,
106     set_score TEXT,
107     game_score TEXT,
108     server TEXT,
109     ending_type TEXT,
110     PRIMARY KEY (video_id, point_number),
111     FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE
112 )
113 """)
114
115 # 4. Create cricket_deliveries table
116 cursor.execute("""
117 CREATE TABLE IF NOT EXISTS cricket_deliveries (
118     video_id TEXT NOT NULL,
119     over_num INTEGER NOT NULL,
120     ball_num INTEGER NOT NULL,
121     start_time REAL NOT NULL,
122     end_time REAL NOT NULL,
123     batsman TEXT,
124     bowler TEXT,
125     runs INTEGER,
126     wicket INTEGER DEFAULT 0,
127     extras INTEGER DEFAULT 0,
128     PRIMARY KEY (video_id, over_num, ball_num),
129     FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE

```

```

130         )
131     """)
132
133     # 5. Create pickleball_rallies table
134     cursor.execute("""
135     CREATE TABLE IF NOT EXISTS pickleball_rallies (
136         video_id TEXT NOT NULL,
137         rally_number INTEGER NOT NULL,
138         start_time REAL NOT NULL,
139         end_time REAL NOT NULL,
140         score_before TEXT,
141         score_after TEXT,
142         shots_count INTEGER,
143         ending_reason TEXT,
144         last_shot_outcome TEXT,
145         PRIMARY KEY (video_id, rally_number),
146         FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE
147     )
148     """)
149
150     # 6. Create tabletennis_rallies table
151     cursor.execute("""
152     CREATE TABLE IF NOT EXISTS tabletennis_rallies (
153         video_id TEXT NOT NULL,
154         rally_number INTEGER NOT NULL,
155         start_time REAL NOT NULL,
156         end_time REAL NOT NULL,
157         score_before TEXT,
158         score_after TEXT,
159         shots_count INTEGER,
160         ending_reason TEXT,
161         last_shot_outcome TEXT,
162         PRIMARY KEY (video_id, rally_number),
163         FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE
164     )
165     """)
166
167     conn.commit()
168     self._migrate(cursor)
169     conn.commit()
170
171     #: The current schema version. Bump when adding a migration step below.
172     SCHEMA_VERSION = 1
173
174     def _migrate(self, cursor) -> None:
175         """Run the versioned migration ladder via PRAGMA user_version.
176
177         Each step is gated on the DB's current version and bumps it after applying.
178         The column-sniffing inside each step stays (so a DB that predates versioning
179         upgrades transparently on first open), but is now scoped and ordered. Future
180         migrations add a new step + bump SCHEMA_VERSION ? no more ad-hoc ALTERs.
181         See docs/CODE_HEALTH_PLAN.md §7 Tier C.
182         """
183         current = cursor.execute("PRAGMA user_version").fetchone()[0]
184         if current >= self.SCHEMA_VERSION:
185             return # up to date
186
187         # --- v1: add last_shot_outcome to rally tables + split video/annotation license
188         if current < 1:
189             # add last_shot_outcome to rally tables that predate the column. CREATE
190             TABLE
191             # IF NOT EXISTS won't alter an existing table, so do it explicitly and
192             # idempotently (skip if the column is already there).
193             for tbl in ("badminton_rallies", "pickleball_rallies",
194 "tabletennis_rallies"):

```

```

193             existing = [r[1] for r in cursor.execute(f"PRAGMA
table_info({tbl})").fetchall()]
194             if "last_shot_outcome" not in existing:
195                 cursor.execute(f"ALTER TABLE {tbl} ADD COLUMN last_shot_outcome
TEXT")
196
197             # split the conflated single license into video vs annotation license.
198             # annotation_license records the dataset's annotation rights (e.g.
ShuttleSet
199             # MIT) for provenance; license_type stays the VIDEO's rights, which gates
200             # commercial use.
201             video_cols = [r[1] for r in cursor.execute("PRAGMA
table_info(videos)").fetchall()]
202             if "annotation_license" not in video_cols:
203                 cursor.execute("ALTER TABLE videos ADD COLUMN annotation_license TEXT")
204
205                 cursor.execute(f"PRAGMA user_version = {self.SCHEMA_VERSION}")
206
207     def insert_video(self, video: VideoMetadata) -> bool:
208         """Insert a video entry, or update it in place if it already exists.
209
210         Uses an UPSERT rather than INSERT OR REPLACE: REPLACE would DELETE the
211         existing row, which (with foreign keys enabled) would cascade-delete the
212         video's child annotations. UPSERT mutates the row in place and leaves
213         children untouched.
214         """
215         with self._get_connection() as conn:
216             cursor = conn.cursor()
217             cursor.execute("""
218             INSERT INTO videos (
219                 video_id, sport, video_url, local_path, license_type,
220                 license_url, annotation_license, is_commercial, status, match_name, fps,
resolution
221                 ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?, ?, ?, ?)
222             ON CONFLICT(video_id) DO UPDATE SET
223                 sport=excluded.sport,
224                 video_url=excluded.video_url,
225                 local_path=excluded.local_path,
226                 license_type=excluded.license_type,
227                 license_url=excluded.license_url,
228                 annotation_license=excluded.annotation_license,
229                 is_commercial=excluded.is_commercial,
230                 status=excluded.status,
231                 match_name=excluded.match_name,
232                 fps=excluded.fps,
233                 resolution=excluded.resolution
234             """, (
235                 video.video_id,
236                 video.sport.value,
237                 video.video_url,
238                 video.local_path,
239                 video.license_type.value,
240                 video.license_url,
241                 video.annotation_license.value if video.annotation_license else None,
242                 1 if video.is_commercial else 0,
243                 video.status.value,
244                 video.match_name,
245                 video.fps,
246                 video.resolution
247             ))
248             conn.commit()
249             return True
250
251     def insert_badminton_rallies(self, rallies: List[BadmintonRally]) -> int:
252         if not rallies:

```

```

253         return 0
254     with self._get_connection() as conn:
255         cursor = conn.cursor()
256         # First clear existing rallies for this video_id to avoid key conflicts on
updates
257         video_id = rallies[0].video_id
258         cursor.execute("DELETE FROM badminton_rallies WHERE video_id = ?",
(video_id,))
259
260         for rally in rallies:
261             cursor.execute("""
262                 INSERT INTO badminton_rallies (
263                     video_id, rally_number, start_time, end_time,
264                     score_before, score_after, shots_count, ending_reason,
last_shot_outcome
265                 ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?)
266             """, (
267                 rally.video_id,
268                 rally.rally_number,
269                 rally.start_time,
270                 rally.end_time,
271                 rally.score_before,
272                 rally.score_after,
273                 rally.shots_count,
274                 rally.ending_reason,
275                 rally.last_shot_outcome
276             ))
277         conn.commit()
278         return len(rallies)
279
280     def insert_tennis_points(self, points: List[TennisPoint]) -> int:
281         if not points:
282             return 0
283         with self._get_connection() as conn:
284             cursor = conn.cursor()
285             video_id = points[0].video_id
286             cursor.execute("DELETE FROM tennis_points WHERE video_id = ?", (video_id,))
287
288             for pt in points:
289                 cursor.execute("""
290                     INSERT INTO tennis_points (
291                         video_id, point_number, start_time, end_time,
292                         set_score, game_score, server, ending_type
293                     ) VALUES (?, ?, ?, ?, ?, ?, ?, ?)
294                 """, (
295                     pt.video_id,
296                     pt.point_number,
297                     pt.start_time,
298                     pt.end_time,
299                     pt.set_score,
300                     pt.game_score,
301                     pt.server,
302                     pt.ending_type
303                 ))
304             conn.commit()
305             return len(points)
306
307     def insert_cricket_deliveries(self, deliveries: List[CricketDelivery]) -> int:
308         if not deliveries:
309             return 0
310         with self._get_connection() as conn:
311             cursor = conn.cursor()
312             video_id = deliveries[0].video_id
313             cursor.execute("DELETE FROM cricket_deliveries WHERE video_id = ?",
(video_id,))

```

```

314
315         for deliv in deliveries:
316             cursor.execute("""
317                 INSERT INTO cricket_deliveries (
318                     video_id, over_num, ball_num, start_time, end_time,
319                     batsman, bowler, runs, wicket, extras
320                 ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?, ?)
321             """, (
322                 deliv.video_id,
323                 deliv.over,
324                 deliv.ball,
325                 deliv.start_time,
326                 deliv.end_time,
327                 deliv.batsman,
328                 deliv.bowler,
329                 deliv.runs,
330                 1 if deliv.wicket else 0,
331                 deliv.extras
332             ))
333         conn.commit()
334         return len(deliveries)
335
336     def insert_pickleball_rallies(self, rallies: List[PickleballRally]) -> int:
337         if not rallies:
338             return 0
339         with self._get_connection() as conn:
340             cursor = conn.cursor()
341             video_id = rallies[0].video_id
342             cursor.execute("DELETE FROM pickleball_rallies WHERE video_id = ?",
343 (video_id,))
344
345         for rally in rallies:
346             cursor.execute("""
347                 INSERT INTO pickleball_rallies (
348                     video_id, rally_number, start_time, end_time,
349                     score_before, score_after, shots_count, ending_reason,
350                     last_shot_outcome
351                 ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?, ?)
352             """, (
353                 rally.video_id,
354                 rally.rally_number,
355                 rally.start_time,
356                 rally.end_time,
357                 rally.score_before,
358                 rally.score_after,
359                 rally.shots_count,
360                 rally.ending_reason,
361                 rally.last_shot_outcome
362             ))
363         conn.commit()
364         return len(rallies)
365
366     def insert_tabletennis_rallies(self, rallies: List[TableTennisRally]) -> int:
367         if not rallies:
368             return 0
369         with self._get_connection() as conn:
370             cursor = conn.cursor()
371             video_id = rallies[0].video_id
372             cursor.execute("DELETE FROM tabletennis_rallies WHERE video_id = ?",
373 (video_id,))
374
375         for rally in rallies:
376             cursor.execute("""
377                 INSERT INTO tabletennis_rallies (
378                     video_id, rally_number, start_time, end_time,

```

```

376         score_before, score_after, shots_count, ending_reason,
last_shot_outcome
377     ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?)
378     """ , (
379         rally.video_id,
380         rally.rally_number,
381         rally.start_time,
382         rally.end_time,
383         rally.score_before,
384         rally.score_after,
385         rally.shots_count,
386         rally.ending_reason,
387         rally.last_shot_outcome
388     ))
389     conn.commit()
390     return len(rallies)
391
392 @staticmethod
393 def _row_to_video(row) -> VideoMetadata:
394     """Rehydrate a DB row into VideoMetadata.
395
396     B5: enum construction (SportType/LicenseType/CurationStatus) used to raise a
397     bare ValueError on any value not in the enum ? a typo or a value from an older
398     schema made the whole query loop explode with a cryptic message and no clue
399     WHICH row was bad. Now: license values unknown to the enum map to UNKNOWN and
400     are logged loudly (OQ2 default: map + log, don't lose data); sport/status have
401     no safe default so they still raise but with the offending video_id in context.
402     """
403     keys = row.keys()
404     ann = row['annotation_license'] if 'annotation_license' in keys else None
405     video_id = row['video_id']
406
407     def _safe_license(value, field: str) -> LicenseType:
408         """Parse a license enum value; map unknown -> UNKNOWN with a warning."""
409         if not value:
410             return LicenseType.UNKNOWN
411         try:
412             return LicenseType(value)
413         except ValueError:
414             logger.warning(
415                 "video %r has unrecognized %s=%r; mapping to UNKNOWN",
416                 video_id, field, value)
417             return LicenseType.UNKNOWN
418
419     try:
420         sport = SportType(row['sport'])
421     except ValueError as e:
422         # No safe default for sport ? surface which row is bad before re-raising.
423         raise ValueError(
424             f"video {video_id!r} has unrecognized sport={row['sport']!r}; "
425             f"cannot rehydrate (fix the DB row or extend SportType)") from e
426
427     try:
428         status = CurationStatus(row['status'])
429     except ValueError as e:
430         raise ValueError(
431             f"video {video_id!r} has unrecognized status={row['status']!r}; "
432             f"cannot rehydrate (fix the DB row)") from e
433
434     return VideoMetadata(
435         video_id=video_id,
436         sport=sport,
437         video_url=row['video_url'],
438         local_path=row['local_path'],
439         license_type=_safe_license(row['license_type'], 'license_type'),

```



```

502         score_after=row['score_after'],
503         shots_count=row['shots_count'],
504         ending_reason=row['ending_reason'],
505         last_shot_outcome=row['last_shot_outcome']
506     ))
507     return rallies
508
509 def get_tennis_points(self, video_id: str) -> List[TennisPoint]:
510     points = []
511     with self._get_connection() as conn:
512         cursor = conn.cursor()
513         cursor.execute("""
514         SELECT * FROM tennis_points
515         WHERE video_id = ?
516         ORDER BY point_number
517         """, (video_id,))
518         for row in cursor.fetchall():
519             points.append(TennisPoint(
520                 video_id=row['video_id'],
521                 point_number=row['point_number'],
522                 start_time=row['start_time'],
523                 end_time=row['end_time'],
524                 set_score=row['set_score'],
525                 game_score=row['game_score'],
526                 server=row['server'],
527                 ending_type=row['ending_type']
528             ))
529     return points
530
531 def get_cricket_deliveries(self, video_id: str) -> List[CricketDelivery]:
532     deliveries = []
533     with self._get_connection() as conn:
534         cursor = conn.cursor()
535         cursor.execute("""
536         SELECT * FROM cricket_deliveries
537         WHERE video_id = ?
538         ORDER BY over_num, ball_num
539         """, (video_id,))
540         for row in cursor.fetchall():
541             deliveries.append(CricketDelivery(
542                 video_id=row['video_id'],
543                 over=row['over_num'],
544                 ball=row['ball_num'],
545                 start_time=row['start_time'],
546                 end_time=row['end_time'],
547                 batsman=row['batsman'],
548                 bowler=row['bowler'],
549                 runs=row['runs'],
550                 wicket=bool(row['wicket']),
551                 extras=row['extras']
552             ))
553     return deliveries
554
555 def get_pickleball_rallies(self, video_id: str) -> List[PickleballRally]:
556     rallies = []
557     with self._get_connection() as conn:
558         cursor = conn.cursor()
559         cursor.execute("""
560         SELECT * FROM pickleball_rallies
561         WHERE video_id = ?
562         ORDER BY rally_number
563         """, (video_id,))
564         for row in cursor.fetchall():
565             rallies.append(PickleballRally(
566                 video_id=row['video_id'],

```

```

567         rally_number=row['rally_number'],
568         start_time=row['start_time'],
569         end_time=row['end_time'],
570         score_before=row['score_before'],
571         score_after=row['score_after'],
572         shots_count=row['shots_count'],
573         ending_reason=row['ending_reason'],
574         last_shot_outcome=row['last_shot_outcome']
575     ))
576     return rallies
577
578 def get_tabletennis_rallies(self, video_id: str) -> List[TableTennisRally]:
579     rallies = []
580     with self._get_connection() as conn:
581         cursor = conn.cursor()
582         cursor.execute("""
583             SELECT * FROM tabletennis_rallies
584             WHERE video_id = ?
585             ORDER BY rally_number
586             """, (video_id,))
587         for row in cursor.fetchall():
588             rallies.append(TableTennisRally(
589                 video_id=row['video_id'],
590                 rally_number=row['rally_number'],
591                 start_time=row['start_time'],
592                 end_time=row['end_time'],
593                 score_before=row['score_before'],
594                 score_after=row['score_after'],
595                 shots_count=row['shots_count'],
596                 ending_reason=row['ending_reason'],
597                 last_shot_outcome=row['last_shot_outcome']
598             ))
599     return rallies
600
601 def get_segment_intervals(self, video_id: str, sport: SportType) -> List[tuple]:
602     """Return chronologically-ordered (start_time, end_time) tuples for a video,
603     regardless of sport. This is the sport-agnostic view used to export
604     indexer-ready evaluation/training labels.
605
606     Sports that have no annotation table yet (e.g. tabletennis, pickleball)
607     yield an empty list rather than raising.
608     """
609     mapping = SEGMENT_TABLES.get(sport)
610     if mapping is None:
611         return []
612     table, order_by = mapping # both come from the trusted map above, never user
input
613     with self._get_connection() as conn:
614         cursor = conn.cursor()
615         cursor.execute(
616             f"SELECT start_time, end_time FROM {table} "
617             f"WHERE video_id = ? ORDER BY {order_by}",
618             (video_id,))
619     )
620     return [(row["start_time"], row["end_time"]) for row in cursor.fetchall()]
621
622 def update_status(self, video_id: str, status: CurationStatus) -> bool:
623     with self._get_connection() as conn:
624         cursor = conn.cursor()
625         cursor.execute("UPDATE videos SET status = ? WHERE video_id = ?",
(status.value, video_id))
626         conn.commit()
627         return cursor.rowcount > 0
628
629 def delete_video(self, video_id: str) -> bool:

```

```

630         """Deletes a video and cascading child entries."""
631         with self._get_connection() as conn:
632             cursor = conn.cursor()
633             # Execute cascade deletes manually just in case SQLite foreign key support
is not active
634             cursor.execute("DELETE FROM badminton_rallies WHERE video_id = ?",
(video_id,))
635             cursor.execute("DELETE FROM tennis_points WHERE video_id = ?", (video_id,))
636             cursor.execute("DELETE FROM cricket_deliveries WHERE video_id = ?",
(video_id,))
637             cursor.execute("DELETE FROM pickleball_rallies WHERE video_id = ?",
(video_id,))
638             cursor.execute("DELETE FROM tabletennis_rallies WHERE video_id = ?",
(video_id,))
639             cursor.execute("DELETE FROM videos WHERE video_id = ?", (video_id,))
640             conn.commit()
641             return cursor.rowcount > 0
642

```

27. user (2026-07-02T04:49:18.258Z)

```

1     from pydantic import BaseModel, Field, model_validator
2     from typing import Optional
3     from enum import Enum
4
5     class SportType(str, Enum):
6         BADMINTON = "badminton"
7         TENNIS = "tennis"
8         TABLE_TENNIS = "tabletennis"
9         PICKLEBALL = "pickleball"
10        CRICKET = "cricket"
11
12    class LicenseType(str, Enum):
13        MIT = "MIT"
14        APACHE_2 = "Apache-2.0"
15        BSD = "BSD"
16        CC0 = "CC0"
17        CC_BY = "CC-BY"
18        CC_BY_NC = "CC-BY-NC"
19        CC_BY_NC_SA = "CC-BY-NC-SA"
20        PUBLIC_DOMAIN = "Public-Domain"
21        PROPRIETARY = "Proprietary"
22        UNKNOWN = "Unknown"
23
24    class CurationStatus(str, Enum):
25        STAGING = "staging"
26        CURATED = "curated"
27        REJECTED = "rejected"
28
29    #: Licenses considered safe for COMMERCIAL use. The single source of truth: B1's
30    #: is_commercial validator uses this directly, and validator.DatasetValidator's
31    #: default whitelist is derived from it ([lt.value for lt in COMMERCIAL_LICENSES],
32    #: via T9). A user can still override via config.yaml's commercial_licenses list.
33    COMMERCIAL_LICENSES = frozenset({
34        LicenseType.MIT, LicenseType.APACHE_2, LicenseType.BSD, LicenseType.CC0,
35        LicenseType.CC_BY, LicenseType.PUBLIC_DOMAIN, LicenseType.PROPRIETARY,
36    })
37
38
39    def is_license_commercial(license_type: LicenseType) -> bool:
40        """True iff `license_type` is in the commercial whitelist (COMMERCIAL_LICENSES)."""
41        return license_type in COMMERCIAL_LICENSES
42
43
44    class VideoMetadata(BaseModel):

```

```

45     video_id: str = Field(..., description="Unique alphanumeric identifier for the
video")
46     sport: SportType
47     video_url: Optional[str] = Field(None, description="Remote URL of the video (if
available)")
48     local_path: Optional[str] = Field(None, description="Local path to video file (if
downloaded or custom)")
49     license_type: LicenseType = Field(
50         ...,
51         description="License of the VIDEO FOOTAGE. This ? not the annotation license ? "
52         "is what gates commercial use (is_commercial). Third-party / broadcast video is
"
53         "'Unknown' unless rights are actually held.",
54     )
55     license_url: Optional[str] = None
56     annotation_license: Optional[LicenseType] = Field(
57         None,
58         description="License of the ANNOTATIONS only (e.g. ShuttleSet's MIT). Recorded
for "
59         "provenance; it does NOT grant any rights to the video footage and is never used
"
60         "for commercial gating.",
61     )
62     is_commercial: bool = Field(False, description="Flag indicating if the VIDEO license
allows commercial use")
63     status: CurationStatus = CurationStatus.STAGING
64     match_name: str = Field(..., description="Name of the match/event")
65     fps: Optional[float] = Field(None, description="Frames per second")
66     resolution: Optional[str] = Field(None, description="Resolution string (e.g. 1080p,
720p)")
67
68     @model_validator(mode="after")
69     def _check_is_commercial_consistency(self) -> "VideoMetadata":
70         """B1/D7: STRICT consistency-check ? is_commercial MUST agree with license_type.
71
72         ``is_commercial`` is an independent bool with no link to ``license_type``, so a
73         caller (or a legacy DB row) could set ``is_commercial=True`` while
74         ``license_type=CC-BY-NC``. The export gate then ships non-commercial footage as
75         "commercially clean" into the indexer's training pool ? a license-compliance
76         landmine for a product whose commercial-clean licensing is non-negotiable.
77
78         This validator makes the invariant un-enforceable-violable: it recomputes the
79         expected flag from ``license_type`` via the whitelist and raises ValueError on
80         any mismatch. Owner-ratified strict (D7, 2026-06-24).
81         """
82         expected = is_license_commercial(self.license_type)
83         if self.is_commercial != expected:
84             raise ValueError(
85                 f"is_commercial={self.is_commercial} is inconsistent with license_type="
86                 f"{self.license_type.value!r} (expected {expected} per the commercial "
87                 f"whitelist). Set is_commercial={expected}, or change license_type."
88             )
89         return self
90

```

28. user (2026-07-02T04:49:19.712Z)

```

1     [build-system]
2     requires = ["setuptools>=61.0.0"]
3     build-backend = "setuptools.build_meta"
4
5     [project]
6     name = "sports-data-collector"
7     version = "0.1.0"
8     description = "A collector and curator for sports video datasets with sport-specific

```

```

schemas and license checks."
9     readme = "README.md"
10    requires-python = ">=3.9"
11    dependencies = [
12        "pydantic>=2.0.0",
13        "pyyaml>=6.0",
14        "tabulate>=0.9.0",
15        # Hardens CVAT XML/zip parsing against entity-expansion / billion-laugh
16        # attacks (see docs/CODE_HEALTH_PLAN.md §7 Tier C). BSD-2-Clause ?
17        # commercially-clean per the project's licensing rule.
18        "defusedxml>=0.7.0",
19    ]
20
21    [project.optional-dependencies]
22    # Video acquisition (`fetch` / crawler). Also requires the external binaries
23    # ffprobe (FFmpeg) and, for YouTube high-res, a JS runtime such as deno ?
24    # neither is pip-installable; see README "External requirements".
25    fetch = [
26        "yt-dlp>=2024.1.0"
27    ]
28    dev = [
29        "pytest>=7.0.0",
30        "pytest-cov>=4.0.0"
31    ]
32    # Per-video observability reports (SOFT dependency: ingestion degrades to a no-op
33    # if absent). Shared, permissively-licensed report toolkit. Pin a tag once the
34    # private repo is published, e.g.:
35    # obsreport @ git+ssh://git@github.com/avidullu/sports-obsreport.git@v0.1.3
36    # For local development install it editable from a sibling checkout:
37    # pip install -e ../sports-obsreport --no-build-isolation
38    report = [
39        "obsreport>=0.1.3",
40    ]
41

```

29. user (2026-07-02T04:49:22.217Z)

```

C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-highlight-
indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\scratchpad\wt-curate-verify\README.md:65:python -m
src.cli.ingest --sport badminton --parser coachai --file-path data/staging_coachai_match.csv
--license CC-BY-NC-SA --video-id staging_match_non_comm
C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-highlight-
indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\scratchpad\wt-curate-verify\src\cli\curate.py:518:
if not video.video_url or "watch?v=example" in video.video_url:
C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-highlight-
indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\scratchpad\wt-curate-verify\src\cli\ingest.py:23:
parser.add_argument("--video-url", default="https://youtube.com/watch?v=example", help="URL to
the video")
C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-highlight-
indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\scratchpad\wt-curate-
verify\src\parsers\badminton\coachai.py:71:         video_url = kwargs.get("video_url") or
"https://youtube.com/watch?v=example"

```

30. assistant (2026-07-02T04:50:31.440Z)

Now let me look at test patterns, config, and check for lint configs.

31. user (2026-07-02T04:50:32.743Z)

```

1     """End-to-end coverage for the SportsDataCLI commands.
2
3     Each test runs against a throwaway config + SQLite DB in tmp_path, so the real
4     data/ directory is never touched. Network (yt-dlp) and the browser are mocked.
5     """

```

```

6     import csv as _csv
7     import json
8     from types import SimpleNamespace
9     from unittest import mock
10
11     import pytest
12     import yaml
13
14     from src.cli.curate import SportsDataCLI
15     from src.core.schemas.base import VideoMetadata, SportType, LicenseType, CurationStatus
16     from src.core.schemas.badminton import BadmintonRally
17
18
19     @pytest.fixture
20     def cli(tmp_path):
21         cfg = {
22             "database_path": str(tmp_path / "test.db"),
23             "video_storage": {
24                 "local_dir": str(tmp_path / "videos"),
25                 "max_download_size_mb": 100,
26                 "auto_download": False,
27                 "max_height": None,
28             },
29             "commercial_licenses": ["MIT", "Apache-2.0", "BSD", "CC0", "CC-BY", "Public-
Domain", "Proprietary"],
30             "supported_sports": ["badminton", "tennis", "tabletennis", "pickleball",
"cricket"],
31         }
32         cfg_path = tmp_path / "config.yaml"
33         cfg_path.write_text(yaml.safe_dump(cfg), encoding="utf-8")
34         return SportsDataCLI(config_path=str(cfg_path))
35
36
37     def _import_args(**kw):
38         base = dict(
39             sport=None, video_path="nonexistent.mp4", annotations_path=None,
40             match_name=None, video_id=None, license="Proprietary", fps=30.0,
resolution="1080p",
41         )
42         base.update(kw)
43         return SimpleNamespace(**base)
44
45
46     def _write_csv(path, header, rows):
47         with open(path, "w", newline="", encoding="utf-8") as f:
48             w = _csv.writer(f)
49             w.writerow(header)
50             for r in rows:
51                 w.writerow(r)
52
53
54     # ----- #
55     # import-local: every sport, both JSON and CSV (covers dispatch + parse) #
56     # ----- #
57
58     def test_import_local_badminton_json(cli, tmp_path):
59         p = tmp_path / "bad.json"
60         p.write_text(json.dumps([
61             {"rally_number": 1, "start_time": 10.0, "end_time": 20.0, "shots_count": 5,
"ending_reason": "winner"},
62             {"rally_number": 2, "start_time": 25.0, "end_time": 30.0, "shots_count": 1,
"ending_reason": "service_fault"},
63         ]), encoding="utf-8")
64         cli.import_local(_import_args(sport="badminton", annotations_path=str(p),
video_id="bad1", match_name="Bad"))

```

```

65         assert len(cli.db.get_badminton_rallies("bad1")) == 2
66
67
68     def test_import_local_pickleball_csv(cli, tmp_path):
69         p = tmp_path / "pb.csv"
70         _write_csv(p, ["rally_number", "start_time", "end_time", "shots_count",
71 "ending_reason"],
72                 [[1, 10.0, 20.0, 6, "winner"], [2, 25.0, 27.0, 1, "service_fault"]])
73         cli.import_local(_import_args(sport="pickleball", annotations_path=str(p),
74 video_id="pb1", match_name="PB"))
75         assert len(cli.db.get_pickleball_rallies("pb1")) == 2
76
77
78     def test_import_local_tabletennis_json(cli, tmp_path):
79         p = tmp_path / "tt.json"
80         p.write_text(json.dumps([
81             {"rally_number": 1, "start_time": 5.0, "end_time": 9.0, "shots_count": 7,
82 "ending_reason": "winner"},
83         ]), encoding="utf-8")
84         cli.import_local(_import_args(sport="tabletennis", annotations_path=str(p),
85 video_id="tt1", match_name="TT"))
86         assert len(cli.db.get_tabletennis_rallies("tt1")) == 1
87
88
89     def test_import_local_tennis_csv(cli, tmp_path):
90         p = tmp_path / "tn.csv"
91         _write_csv(p, ["point_number", "start_time", "end_time", "server", "ending_type"],
92                 [[1, 10.0, 20.0, "player1", "winner"]])
93         cli.import_local(_import_args(sport="tennis", annotations_path=str(p),
94 video_id="tn1", match_name="TN"))
95         assert len(cli.db.get_tennis_points("tn1")) == 1
96
97
98     def test_import_local_cricket_json(cli, tmp_path):
99         p = tmp_path / "ck.json"
100         p.write_text(json.dumps([
101             {"over": 0, "ball": 1, "start_time": 5.0, "end_time": 15.0, "runs": 4, "wicket":
102 False},
103         ]), encoding="utf-8")
104         cli.import_local(_import_args(sport="cricket", annotations_path=str(p),
105 video_id="ck1", match_name="CK"))
106         assert len(cli.db.get_cricket_deliveries("ck1")) == 1
107
108
109     # ----- #
110     # import-local: error / edge branches #
111     # ----- #
112
113     def test_import_local_unsupported_sport_exits(cli, tmp_path):
114         p = tmp_path / "x.json"
115         p.write_text("[]", encoding="utf-8")
116         with pytest.raises(SystemExit):
117             cli.import_local(_import_args(sport="football", annotations_path=str(p)))
118
119
120     def test_import_local_missing_annotations_exits(cli):
121         with pytest.raises(SystemExit):
122             cli.import_local(_import_args(sport="badminton",
123 annotations_path="does_not_exist.json"))
124
125
126     def test_import_local_overlap_fails_validation(cli, tmp_path):
127         p = tmp_path / "ov.csv"
128         _write_csv(p, ["rally_number", "start_time", "end_time"],

```

Sports Data Collector Configuration

Path to the SQLite database

```
database_path: "./data/sports_metadata.db"
```

Storage configuration for downloaded videos

```
video_storage:
  # Folder where downloaded or local videos will be stored
  local_dir: "./data/videos"
  # Maximum allowed size for automatic downloads (in Megabytes)
  max_download_size_mb: 12000
  # Whether to attempt video downloads automatically during ingestion
  auto_download: true
  # Optional cap on download resolution height (e.g. 1080 to avoid pulling 4K).
  # Leave unset/null to always fetch the best resolution available.
  max_height: null
```

Whitelist of licenses considered safe for commercial use

```
commercial_licenses:
  - "MIT"
  - "Apache-2.0"
  - "BSD"
  - "CC0"
  - "CC-BY"
  - "Public-Domain"
  - "Proprietary" # Used for custom user-added local videos
```

Ingestion limits

```
ingestion_limits:
  max_total_download_size_mb: 60000
  max_total_video_count: 10
```

Whitelist of all sports supported in the system

```
supported_sports:
  - "badminton"
  - "tennis"
  - "tabletennis"
  - "pickleball"
  - "cricket"
```

Validation / parsing thresholds. Each has a rationale comment; override here only with a deliberate reason. See docs/CODE_HEALTH_PLAN.md §7 Tier C.

```
validation:
  coachai:
    # Seconds subtracted from the first stroke (serve) to get the rally start.
    # 1.0s covers the serve wind-up without eating the previous rally's tail.
    start_buffer: 1.0
    # Seconds added to the last stroke to get the rally end. 2.0s covers the
    # shuttle settling / umpire call after the last shot.
    end_buffer: 2.0
  rally_cvat:
    # Rallies with fewer than this many shots/second are flagged "shot_density"
    # (medium) ? the span may be inflated by a stray box. 0.3/s ? a shot every
    # ~3.3s, well below any real rally tempo.
    min_shots_per_second: 0.3
    # A gap between consecutive rallies longer than this is flagged
    # "possible_missed_rally" (medium). 15.0s catches an omitted rally that left
```



```

# a visible hole without false-firing on normal between-rally breaks.
max_gap_s: 15.0
# The ±boundary bar (seconds) rally first/last boxes must meet. 0.3s is the
# golden-quality tolerance from Brief §4; a coarser sampling grid makes it
# structurally unmeetable (the boundary_grid_snap blocker).
boundary_tolerance_s: 0.3
ci.yml
name: CI

on:
  push:
    branches: [master]
  pull_request:

jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-python@v5
        with:
          python-version: "3.11"
      - name: Install dependencies
        run: |
          python -m pip install --upgrade pip
          pip install -e ".[dev]"
      - name: Run tests
        run: python -m pytest -q
.
..
.git
.gitattributes
.github
.gitignore
README.md
cleanup.py
config
data
C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/tests/test_db.py

```

33. assistant (2026-07-02T04:50:42.358Z)

CI is pytest-only; tests use tmp_path + throwaway config. Now let me verify the two G: candidate files against DB metadata and the golden label CSVs.

34. user (2026-07-02T04:50:47.014Z)

```

{
  "programs": [

  ],
  "stream_groups": [

  ],
  "streams": [
    {
      "width": 1920,
      "height": 1080,
      "r_frame_rate": "30000/1001",
      "avg_frame_rate": "30000/1001",
      "nb_frames": "13418"
    }
  ],
  "format": {

```

```

        "duration": "447.713933",
        "size": "196182380"
    }
}

```

35. user (2026-07-02T04:50:47.309Z)

```

29,359.664,366.661,winner,badminton
30,374.162,381.659,winner,badminton
31,394.661,399.412,winner,badminton
===
94,1370.997,1378.001,unforced_error,badminton
95,1385.257,1391.510,forced_error,badminton
96,1398.257,1405.757,unforced_error,badminton
===
rally_number,start_time,end_time,ending_reason,sport
1,37.434,72.194,winner,badminton

```

36. user (2026-07-02T04:51:04.842Z)

```

{
  "programs": [

  ],
  "stream_groups": [

  ],
  "streams": [
    {
      "width": 1920,
      "height": 1080,
      "r_frame_rate": "60000/1001",
      "avg_frame_rate": "60000/1001",
      "nb_frames": "84480"
    }
  ],
  "format": {
    "duration": "1409.408000",
    "size": "11445052772"
  }
}

```

37. user (2026-07-02T04:51:52.680Z)

Normalize line endings: store text as LF in the repo, check out native.

```
* text=auto
```

Treat data/media as binary ? never line-ending-convert or text-diff these.

```

*.db binary
*.sqlite binary
*.sqlite3 binary
*.mp4 binary
*.png binary
*.jpg binary
*.jpeg binary
staging_match_non_comm rallies: 2
shuttlest_12 rallies: 79

```

38. user (2026-07-02T04:52:19.918Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/src/cli/curate.py has been updated successfully. (file state is current in your context ? no need to Read it back)

39. user (2026-07-02T04:52:43.201Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/src/cli/curate.py has been updated successfully. (file state is current in your context ? no need to Read it back)

40. user (2026-07-02T04:52:52.654Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/src/cli/curate.py has been updated successfully. (file state is current in your context ? no need to Read it back)

41. user (2026-07-02T04:52:59.483Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/src/cli/curate.py has been updated successfully. (file state is current in your context ? no need to Read it back)

42. assistant (2026-07-02T04:53:04.160Z)

Now the tests. Following the `test\_cli.py` fixture pattern, in a new `tests/test\_verify.py`.

43. assistant (2026-07-02T04:53:29.987Z)

```
"""Tests for `curate verify` ? the registry-vs-disk local_path audit.
```

```
Every test runs against a throwaway config + SQLite DB in tmp_path (same pattern as test_cli.py), so the real data/ directory is never touched.
```

```
"""
```

```
from types import SimpleNamespace
```

```
import pytest
import yaml
```

```
from src.cli.curate import CurationError, SportsDataCLI
from src.core.schemas.base import CurationStatus, LicenseType, SportType, VideoMetadata
```

```
@pytest.fixture
```

```
def cli(tmp_path):
```

```
    cfg = {
```

```
        "database_path": str(tmp_path / "test.db"),
```

```
        "video_storage": {
```

```
            "local_dir": str(tmp_path / "videos"),
```

```
            "max_download_size_mb": 100,
```

```
            "auto_download": False,
```

```
            "max_height": None,
```

```
        },
```

```
        "commercial_licenses": ["MIT", "Apache-2.0", "BSD", "CC0", "CC-BY", "Public-Domain",
```

```
        "Proprietary"],
```

```
        "supported_sports": ["badminton", "tennis", "tabletennis", "pickleball", "cricket"],
```

```
    }
```

```
    cfg_path = tmp_path / "config.yaml"
```

```
    cfg_path.write_text(yaml.safe_dump(cfg), encoding="utf-8")
```

```
    return SportsDataCLI(config_path=str(cfg_path))
```

```
def _video(video_id, *, local_path=None, video_url=None, sport=SportType.BADMINTON,
```

```
          status=CurationStatus.CURATED, license_type=LicenseType.PROPRIETARY):
```

```
    return VideoMetadata(
```

```
        video_id=video_id, sport=sport, video_url=video_url, local_path=local_path,
```

```
        license_type=license_type, license_url=None,
```

```
        is_commercial=(license_type == LicenseType.PROPRIETARY),
```

```
        status=status, match_name=video_id,
```

```

)

def _verify_args(**kw):
    base = dict(sport=None, status="all", strict=False)
    base.update(kw)
    return SimpleNamespace(**base)

----- #
verify_local_paths: classification #
----- #

def test_verify_classifies_ok_missing_and_pathless(cli, tmp_path):
    real = tmp_path / "real.mp4"
    real.write_bytes(b"\x00" * 16)
    cli.db.insert_video(_video("v_ok", local_path=str(real)))
    cli.db.insert_video(_video("v_gone", local_path=str(tmp_path / "moved_away.mp4")))
    cli.db.insert_video(_video("v_urlonly", video_url="https://www.youtube.com/watch?v=abc",
                               license_type=LicenseType.UNKNOWN))

    report = cli.verify_local_paths()

    assert [v.video_id for v in report.ok] == ["v_ok"]
    assert [v.video_id for v in report.missing] == ["v_gone"]
    assert [v.video_id for v in report.pathless] == ["v_urlonly"]
    assert report.n_videos == 3

def test_verify_covers_every_status_exactly_once(cli, tmp_path):
    gone = str(tmp_path / "gone.mp4")
    cli.db.insert_video(_video("v_staging", local_path=gone, status=CurationStatus.STAGING,
                               license_type=LicenseType.UNKNOWN))
    cli.db.insert_video(_video("v_curated", local_path=gone, status=CurationStatus.CURATED))
    cli.db.insert_video(_video("v_rejected", local_path=gone, status=CurationStatus.REJECTED,
                               license_type=LicenseType.UNKNOWN))

    report = cli.verify_local_paths()
    assert sorted(v.video_id for v in report.missing) == ["v_curated", "v_rejected",
"v_staging"]
    assert report.n_videos == 3 # one bucket entry per row, no dupes

    only_curated = cli.verify_local_paths(statuses=[CurationStatus.CURATED])
    assert [v.video_id for v in only_curated.missing] == ["v_curated"]

def test_verify_relative_path_resolves_against_cwd(cli, tmp_path, monkeypatch):
    (tmp_path / "videos").mkdir(exist_ok=True)
    (tmp_path / "videos" / "rel.mp4").write_bytes(b"\x00")
    cli.db.insert_video(_video("v_rel", local_path="./videos/rel.mp4"))

    monkeypatch.chdir(tmp_path)
    assert [v.video_id for v in cli.verify_local_paths().ok] == ["v_rel"]

    # From anywhere else the same relative path is dangling ? same semantics as
    # every other local_path consumer (export / play harness).
    monkeypatch.chdir(tmp_path / "videos")
    assert [v.video_id for v in cli.verify_local_paths().missing] == ["v_rel"]

def test_verify_sport_filter_and_unknown_sport(cli, tmp_path):
    cli.db.insert_video(_video("v_bad", local_path=str(tmp_path / "gone.mp4")))
    cli.db.insert_video(_video("v_tt", sport=SportType.TABLE_TENNIS,

```

```

        local_path=str(tmp_path / "gone2.mp4"))

report = cli.verify_local_paths(sport="badminton")
assert [v.video_id for v in report.missing] == ["v_bad"]

with pytest.raises(CurationError):
    cli.verify_local_paths(sport="football")

```

----- # **CLI wrapper: output + exit codes** # -----

```

def test_verify_cli_reports_and_exits_zero_without_strict(cli, tmp_path, capsys):
    cli.db.insert_video(_video("v_gone", local_path=str(tmp_path / "gone.mp4")))

    cli.verify(_verify_args()) # must NOT raise SystemExit without --strict

    out = capsys.readouterr().out
    assert "MISSING : 1" in out
    assert "v_gone" in out
    assert "Dangling rows" in out

def test_verify_cli_strict_exits_nonzero_on_dangling(cli, tmp_path):
    cli.db.insert_video(_video("v_gone", local_path=str(tmp_path / "gone.mp4")))

    with pytest.raises(SystemExit) as exc:
        cli.verify(_verify_args(strict=True))
    assert exc.value.code == 1

def test_verify_cli_strict_passes_when_all_resolve(cli, tmp_path, capsys):
    real = tmp_path / "real.mp4"
    real.write_bytes(b"\x00")
    cli.db.insert_video(_video("v_ok", local_path=str(real)))

    cli.verify(_verify_args(strict=True)) # no dangling rows -> no SystemExit

    out = capsys.readouterr().out
    assert "Resolved OK : 1" in out
    assert "MISSING : 0" in out

def test_verify_cli_status_filter(cli, tmp_path, capsys):
    cli.db.insert_video(_video("v_staging_gone", local_path=str(tmp_path / "gone.mp4"),
                               status=CurationStatus.STAGING, license_type=LicenseType.UNKNOWN))

    # Only curated rows requested -> the staging dangler is out of scope.
    cli.verify(_verify_args(status="curated", strict=True))
    assert "No matching videos registered" in capsys.readouterr().out

def test_verify_cli_flags_placeholder_url_rows(cli, capsys):
    cli.db.insert_video(_video("v_placeholder", video_url="https://youtube.com/watch?v=example",
                               status=CurationStatus.STAGING, license_type=LicenseType.UNKNOWN))
    cli.db.insert_video(_video("v_urlonly", video_url="https://www.youtube.com/watch?v=real123",
                               status=CurationStatus.STAGING, license_type=LicenseType.UNKNOWN))

    cli.verify(_verify_args())

    out = capsys.readouterr().out
    assert "no local_path and no real source URL" in out # the placeholder row

```

```
assert "no local_path (URL-only entry)" in out          # the real-URL row
```

44. user (2026-07-02T04:53:30.063Z)

File created successfully at: C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/tests/test_verify.py (file state is current in your context ? no need to Read it back)

45. assistant (2026-07-02T04:53:38.863Z)

Now run the new tests plus the existing CLI tests.

46. user (2026-07-02T04:53:44.559Z)

```
..... [100%]  
40 passed in 2.41s
```

47. assistant (2026-07-02T04:53:49.583Z)

All pass. Now document the command in the README.

48. user (2026-07-02T04:53:55.546Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

49. user (2026-07-02T04:54:03.356Z)

```
96     ### 6. Export Annotations as Indexer-Ready Eval/Training Sets  
97     To turn curated annotations into ground-truth files the badminton-highlight-indexer can  
98     score and train against, export per-video `start,end` CSVs plus a `manifest.json`:  
99     ```bash  
100     # Curated + commercially-safe videos across all supported sports (default):  
101     python -m src.cli.curate export --out-dir data/eval_export  
102     # One sport, also including staging entries:  
103     python -m src.cli.curate export --sport badminton --include-staging  
104     ```  
105     The exporter writes `//.csv>` (decimal-second `start,end` rows)  
106     and a `manifest.json` pairing each CSV with its video `local_path`/`video_url`, fps, resolution,  
107     and license. The indexer keys videos by **file MD5**, not by our `video_id`, so the downstream  
108     eval/training run must process the exact video file named in the manifest. Flags: `--sport`  
109     (restrict to one sport), `--include-staging` (default: curated only), `--include-noncommercial`  
110     (default: commercial-safe only).  
111     ---  
112     ## Configuration  
113     Adjust settings in `config/config.yaml`:  
114     - **`database_path`**: Path to the SQLite database.  
115     - **`video_storage`**: Local folder path and automatic download size limit.  
116     - **`commercial_licenses`**: Whitelist of licenses allowed for commercial use.  
117     - **`supported_sports`**: Allowed sports categories.  
118     ---  
119     ## Documentation  
120
```

50. user (2026-07-02T04:54:17.913Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify/README.md has been updated successfully. (file state is current in your context ? no need to Read it back)

51. assistant (2026-07-02T04:54:35.714Z)

Now the DB surgery on the worktree copy. I'll use the repo's own `SportsDatabase.delete\_video` for the cascade delete, plus targeted UPDATES.

52. user (2026-07-02T04:54:46.710Z)

re-pointed adarsh_avi_singles -> G:\My Drive\Sharing\Badminton\Boxhill\June 1 2026\Adarsh and Avi.MP4
re-pointed mahadevpura_singles -> G:\My Drive\Sharing\Golden Annotations\MahadevpuraSingles.MP4
shuttleset_12: placeholder video_url cleared (row + 79 rallies kept)
staging_match_non_comm: purged (demo placeholder row)
('shuttleset_12', None, None)
('adarsh_avi_singles', None, 'G:\My Drive\Sharing\Badminton\Boxhill\June 1 2026\Adarsh and Avi.MP4')
('mahadevpura_singles', None, 'G:\My Drive\Sharing\Golden Annotations\MahadevpuraSingles.MP4')
total videos: 30
shuttleset_12 rallies: 79

53. user (2026-07-02T04:54:59.651Z)

Exit code 1
Traceback (most recent call last):
File "<string>", line 7, in <module>
cli.verify(SimpleNamespace(sport=None, status='all', strict=False))
^^^^^^^^^^
AttributeError: 'SportsDataCLI' object has no attribute 'verify'

54. assistant (2026-07-02T04:55:08.687Z)

The main checkout's `src` (old code) shadowed the worktree on sys.path. Fix the import order.

55. user (2026-07-02T04:55:12.995Z)

=== Registry local_path verification ===
+-----+-----+-----+-----+
+-----+
| Video ID | Status | File | Local path |
|
+=====+
+-----+
| adarsh_avi_singles | curated | OK | G:\My Drive\Sharing\Badminton\Boxhill\June 1 2026\Adarsh and Avi.MP4 |
+-----+
+-----+
| gbaaddy | curated | MISSING | C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\gBaaddy20May2026.MP4 |
+-----+
+-----+
| kushagra_singles | curated | MISSING | C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\KushagraYashAviFirstVideo.MP4 |
+-----+
+-----+
| largetest_doubles | curated | MISSING | C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\LargeTestVideo.MP4 |
+-----+
+-----+
| mahadevpura_singles | curated | OK | G:\My Drive\Sharing\Golden Annotations\MahadevpuraSingles.MP4 |
+-----+
+-----+
| my_custom_local_match | curated | MISSING | data/videos/my_custom_match.mp4 |
|

shuttleset_1	staging	OK	./data/videos\shuttleset_1.mp4
shuttleset_10	staging	OK	./data/videos\shuttleset_10.mp4
shuttleset_11	staging	OK	./data/videos\shuttleset_11.mp4
shuttleset_12	staging	-	no local_path and no real source URL
shuttleset_13	staging	OK	./data/videos\shuttleset_13.mp4
shuttleset_14	staging	OK	./data/videos\shuttleset_14.mp4
shuttleset_15	staging	OK	./data/videos\shuttleset_15.mp4
shuttleset_16	staging	OK	./data/videos\shuttleset_16.mp4
shuttleset_17	staging	OK	./data/videos\shuttleset_17.mp4
shuttleset_18	staging	OK	./data/videos\shuttleset_18.mp4
shuttleset_19	staging	OK	./data/videos\shuttleset_19.mp4
shuttleset_2	staging	OK	./data/videos\shuttleset_2.mp4
shuttleset_20	staging	OK	./data/videos\shuttleset_20.mp4
shuttleset_21	curated	OK	./data/videos\shuttleset_21.mp4
shuttleset_22	staging	OK	./data/videos\shuttleset_22.mp4
shuttleset_23	staging	OK	./data/videos\shuttleset_23.mp4


```

-----+
| shuttleset_24      | staging | OK      | ./data/videos\shuttleset_24.mp4
|
+-----+-----+-----+-----+
-----+
| shuttleset_25      | staging | OK      | ./data/videos\shuttleset_25.mp4
|
+-----+-----+-----+-----+
-----+
| shuttleset_3       | staging | OK      | ./data/videos\shuttleset_3.mp4
|
+-----+-----+-----+-----+
-----+
| shuttleset_4       | staging | OK      | ./data/videos\shuttleset_4.mp4
|
+-----+-----+-----+-----+
-----+
| shuttleset_6       | staging | OK      | ./data/videos\shuttleset_6.mp4
|
+-----+-----+-----+-----+
-----+
| shuttleset_8       | staging | OK      | ./data/videos\shuttleset_8.mp4
|
+-----+-----+-----+-----+
-----+
| shuttleset_9       | staging | OK      | ./data/videos\shuttleset_9.mp4
|
+-----+-----+-----+-----+
-----+
| testlarge_short    | curated | MISSING | C:\Users\avidu\Projects\Annotation
Setup\Collect\Golden Labelled\TestLargeVideo.MP4
|
+-----+-----+-----+-----+
-----+

```

Resolved OK : 24 MISSING : 5 No local_path : 1

Dangling rows (registered local_path does not resolve on disk):

- my_custom_local_match: data/videos/my_custom_match.mp4
- kushagra_singles: C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\KushagraYashAviFirstVideo.MP4
- targetest_doubles: C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\LargeTestVideo.MP4
- gbaaddy: C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\gBaaddy20May2026.MP4
- testlarge_short: C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\TestLargeVideo.MP4

Re-point these rows at the real file location (or re-`fetch` the video); until then `export` emits md5=None for them, breaking the indexer join.

56. assistant (2026-07-02T04:55:23.117Z)

The command works end-to-end. Now lint the changed files.

57. user (2026-07-02T04:55:26.486Z)

ruff 0.15.16
All checks passed!

58. assistant (2026-07-02T04:55:41.238Z)

Ruff clean; no mypy config exists in this repo (pyproject has no tool sections). Now commit ? verifying branch state first per the safety protocol.

59. user (2026-07-02T04:55:43.043Z)

```
fix/curate-verify-registry-rot
adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)
  M README.md
  M data/sports_metadata.db
  M src/cli/curate.py
?? tests/test_verify.py
```

60. user (2026-07-02T04:55:55.937Z)

```
warning: in the working copy of 'tests/test_verify.py', LF will be replaced by CRLF the next
time Git touches it
[fix/curate-verify-registry-rot aef4621] feat(curate): add `verify` local_path audit; fix
golden-corpus registry rot
 4 files changed, 297 insertions(+)
 create mode 100644 tests/test_verify.py
aef4621 feat(curate): add `verify` local_path audit; fix golden-corpus registry rot
```

61. assistant (2026-07-02T04:56:02.107Z)

Only my commit on top of origin/master. Push and open the PR.

62. user (2026-07-02T04:56:07.071Z)

```
remote:
remote: Create a pull request for 'fix/curate-verify-registry-rot' on GitHub by visiting:
remote:      https://github.com/Khelsutra/sports-data-collector/pull/new/fix/curate-verify-
registry-rot
remote:
branch 'fix/curate-verify-registry-rot' set up to track 'origin/fix/curate-verify-registry-rot'.
To https://github.com/Khelsutra/sports-data-collector.git
 * [new branch]      fix/curate-verify-registry-rot -> fix/curate-verify-registry-rot
```

63. user (2026-07-02T04:56:44.280Z)

<https://github.com/Khelsutra/sports-data-collector/pull/55>

64. user (2026-07-02T04:57:13.190Z)

Structured output provided successfully