

claude-windows / c1fcb4aa-537d-41fd-858e-53f93349bf5a

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\subagents\workflows\wf_7049e1ed-b93\agent-a5438b39754d34cc6.jsonl`  
- SHA-256: `cb1a5d17d323688f23348bc392cb350d55f6c8d4694fa73b226dbd5c29215f06`  
- Source modified: `2026-07-02T04:58:56+00:00`  
- Imported at: `2026-07-05T16:48:24+00:00`  
- project: `wf_7049e1ed-b93`  
- session_id: `c1fcb4aa-537d-41fd-858e-53f93349bf5a`
```

Transcript

1. user (2026-07-02T04:57:13.270Z)

You are an adversarial code reviewer. Today is 2026-07-02. Review PR #55 in the repo at C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector (gh pr view 55, gh pr diff 55; the working copy is in worktree C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify ? read code there for context).

The PR claims to fix: R2-6 (golden-corpus registry rot) + critic #1 (evidence in C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/AUDIT_REPORT_khelsutra-guru_2026-07-02.md ? Grep for the IDs; honor the verifier notes' scoping corrections).

Hunt REAL defects: correctness bugs in the diff, broken/vacuous tests, behavior changes beyond the brief, staging accidents (files that should not be in the diff ? check the file list!), convention violations (default-ON when brief says OFF, missing schema-version check, CWD assumptions), and whether the fix actually addresses the finding. READ-ONLY ? do not modify or comment on GitHub. blocker = must not merge; major = fix before merge; minor = nice-to-have.

2. user (2026-07-02T04:57:20.181Z)

```
title: feat(curate): add `verify` local_path audit; fix golden-corpus registry rot  
state: OPEN  
author: avidullu (Avi)  
labels:  
assignees:  
reviewers:  
projects:  
milestone:  
number: 55  
url: https://github.com/Khelsutra/sports-data-collector/pull/55  
additions: 297  
deletions: 0  
auto-merge: disabled  
--
```

Fixes audit finding **R2-6** (golden-corpus registry rot) + critic finding #1 (placeholder junk rows) from the 2026-07-02 workspace audit.

What / why

All 6 owned 'curated' golden rows in `data/sports\_metadata.db` pointed at `C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\\*.MP4` ? a directory that no longer exists (the folder moved to `khelsutra-guru/Annotation Setup/`, which now holds only the `rallies.csv` label files). Nothing in the toolchain re-checks `local\_path` after registration, so `curate export` silently emits `md5=None` for the whole owned corpus, breaking the indexer's deterministic MD5 join (PR #51 contract).

1. ****New `curate verify` subcommand**** ? resolves every registered `videos.local\_path` against the disk and reports dangling rows (plus rows with no `local\_path` at all, noting when they also lack a real source URL). `--strict` exits non-zero when any dangling row exists so it can gate CI/cron/exports; `--sport` / `--status` filter. Library core `verify\_local\_paths()` returns a structured `VerifyReport` (same pattern as `import\_annotations`). README documents it (§7).
2. ****Re-pointed 2 of the 6 owned golden rows**** to copies I personally verified on disk:
 - `adarsh\_avi\_singles` ? `G:\My Drive\Sharing\Badminton\Boxhill\June 1 2026\Adarsh and Avi.MP4` ? exists (11,445,052,772 B); ffprobe: 1920x1080 @ 59.94 fps (matches the DB row), duration 1409.408 s vs last golden label ending 1405.757 s (labels span the whole file).
 - `mahadevpura\_singles` ? `G:\My Drive\Sharing\Golden Annotations\MahadevpuraSingles.MP4` ? exists (196,182,380 B, matching the 187 MB documented for the F: original); ffprobe: 1920x1080 @ 29.97 fps (matches the DB row), duration 447.714 s vs last label 399.412 s.
 - Note: these are Google Drive File Stream paths (content streams on demand); no persisted MD5s of the originals exist anywhere to compare against, so identity rests on filename + fps + resolution + label-duration consistency, all of which match exactly.
3. ****Left 4 rows dangling on purpose**** ? `kushagra\_singles`, `largetest\_doubles`, `gbaaddy`, `testlarge\_short` exist nowhere reachable today. Their only documented copy is `F:\[Khelsutra] GoPro Hero Black 12\KhelSutra\Training\Golden Labelled\` (badminton-highlight-indexer `docs/DATA\_IN\_GCS.md` §5/§6), and the F: drive is ****not mounted**** (only C:/G: exist; verified via `Get-PSDrive` + recursive filename search of `C:\Users\avidu\Projects` and `G:\My Drive`, and `gs://khelsutra-rally-corpus/videos/` holds none of them per the audit). Re-pointing them at an unverifiable path would defeat the point of `verify`; they stay flagged until F: is mounted (then: verify, hash, re-point ? and ideally execute the DATA_IN_GCS §6 Phase-2 proxy+upload plan).
4. ****Placeholder junk rows**** (critic finding):
 - `staging\_match\_non\_comm` ? purged. Pure demo fixture from the README ingest example (fake `watch?v=example` URL, 2 rallies derived from the committed `data/staging\_coachai\_match.csv`); fully recreatable by re-running the documented command.
 - `shuttleset\_12` ? kept (it holds 79 real ShuttleSet rally annotations) but its fake `watch?v=example` URL is cleared to NULL; `fetch` already treats a NULL URL as "no real source URL" identically, and `verify` now reports it as "no local_path and no real source URL".

Binary-diff note

`data/sports\_metadata.db` is a git-tracked binary sqlite (marked `binary` in `.gitattributes`) ? this PR rewrites the file, so GitHub shows no textual diff. Contents changed: 2 `local\_path` UPDATES, 1 `video\_url` ? NULL, 1 row purged (+ its 2 `badminton\_rallies` rows via the repo's own `delete\_video` cascade), then `VACUUM`. Row count 31 ? 30. The committed `data/eval\_export/` artifact is ****not**** tracked in git, so it is untouched here; re-run `curate export` after merge to refresh the manifest (it will now emit real MD5s for the two re-pointed videos, at Drive-streaming speed for the 10.7 GiB one).

Verification

Ran locally (Windows, Python 3.13):

- `python -m pytest tests/test\_verify.py tests/test\_cli.py -q` ? 40 passed (9 new tests: classification ok/missing/pathless, per-status coverage, relative-path-vs-cwd semantics, sport filter + unknown sport, `--strict` exit codes both ways, status filter, placeholder-URL detail strings; all on a tmp_path DB).
- `ruff check src/cli/curate.py tests/test\_verify.py` ? clean (repo has no ruff/mypy config; CI is pytest-only).
- End-to-end `curate verify` against the updated committed DB from the real checkout: 24 OK / 5 MISSING / 1 no-local-path ? the 5 MISSING are the 4 F:-only golden rows above plus `my\_custom\_local\_match` (a pre-existing README demo row, out of scope, left as-is).
- Both re-pointed files verified present via `ls`/ffprobe immediately before the DB update (the update script asserts `os.path.exists` and refuses otherwise).

CI covers the full pytest suite on ubuntu (the new tests are OS-independent: tmp_path fixtures only).

? Generated with [Claude Code](https://claude.com/claude-code)

3. user (2026-07-02T04:57:21.563Z)

```

476-- verdict: CONFIRMED
477-[Omitted long context line]
478-[Omitted long context line]
479-
480-### R2-5 [P1/S] ? badminton-highlight-indexer
481-**Engine CLAUDE.md 'Current state' section is ~3 weeks false ? claims platform/owned-model
work 'has not started'**
482-- verdict: CONFIRMED
483-[Omitted long context line]
484-[Omitted long context line]
485-
486-### R2-6 [P1/M] ? sports-data-collector
487-**Golden-corpus registry rot: all 6 owned 'curated' videos point at a deleted absolute
local_path and the 4.9GB source videos are single-copy with no documented backup**
488-- verdict: CONFIRMED
489-[Omitted long context line]
490-[Omitted long context line]
491-
492-### R2-7 [P2/S] ? non-git
493-**Roora vendor pilot record (paid deliverables + corrected golden labels + 3 videos) is a
single unversioned local copy at risk of loss**
494-- verdict: CONFIRMED
495-[Omitted long context line]
496-[Omitted long context line]

```

4. user (2026-07-02T04:57:30.154Z)

```

README.md
data/sports_metadata.db
src/cli/curate.py
tests/test_verify.py
---
aef46218a45c3797b15564d580bbb240e67caa30 feat(curate): add `verify` local_path audit; fix
golden-corpus regist?

```

5. user (2026-07-02T04:57:30.683Z)

```

<persisted-output>
Output too large (40.4KB). Full output saved to: C:\Users\avidu\.claude\projects\C--Users-avidu-
Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\tool-
results\b23gfq37j.txt

```

Preview (first 2KB):

Correctly-scoped fix (S): (a) don't start the 333 ms clock until the panel is first shown and pause/clear it when hidden ? gate it inside `toggle()` in panel.ts; (b) make render's clock read non-scanning (e.g. a `peekNow()` on DirectVideoHandler that uses only the cached element), keeping the full scan for explicit triggers (Refresh button, panel open). Optional follow-up (M, separate decision): replace the `<all\_urls>` content script with activeTab + scripting.executeScript on icon click, which eliminates injection on every page entirely and drops the broad host permission ? aligns with LD-7's own rationale, worth raising but not required to fix the perf debt. Not OWNER_GATED ? pure code change, no spend or product pivot.

R2-2 [P2/S] ? rally-annotator

```

**Locked decision LD-i8 (web language persisted + seeded from browser language) is documented as
shipped but not implemented ? panel always boots in English and forgets the choice on reload**
- verdict: CONFIRMED
- reasoning: Doc claims shipped: rally-annotator/docs/LOCALIZATION.md:55-57 (LD-i8: "Web:
persisted in chrome.storage.local (seeded from the browser language; default en)") and :99-100
mark the web i18n phases done (PR #18 "selector", PR #19); web/README.md:22-23 says "persisted
in-panel language selector". Code says otherwise: (1) web/src/i18n/index.ts:56 boots the module-
level locale to DEFAULT_LOCALE ("en") ? a fresh content-script injection always starts English;
(2) negotiateLocale (index.ts:72-80), which implements the stored?navigator.languages?en order,
has ZERO runtime callers ? grep across web/ shows its only callers are unit tests
(web/test/i18n.test.ts:47-52); (3) panel.ts:18 declares onLocaleChange with the comment

```

"optional: persist the choice (content script wires this)", but the actual mountPanel call in web/entrypoints/content.ts:32-42 passes only annotator/video/identity/download ? onLocaleChange is never wired, so the change handler at panel.ts:306-311 updates only the in-memory `curre
...
</persisted-output>

6. user (2026-07-02T04:57:38.676Z)

27:### R2-6 [P1/M] ? sports-data-collector
28-**Golden-corpus registry rot: all 6 owned 'curated' videos point at a deleted absolute local_path and the 4.9GB source videos are single-copy with no documented backup**
29-- verdict: CONFIRMED
30-- reasoning: Registry rot half CONFIRMED by direct observation. sports-data-collector/data/sports_metadata.db (git-tracked; rows landed in commit 9e6f204 "#19 ingest the n=6 owned golden corpus", DB last touched 9f1e792 with no path fix since): videos table has 8 status='curated' rows; the 6 owned ones (adarsh_avi_singles, kushagra_singles, largetest_doubles, gbaaddy, testlarge_short, mahadevpura_singles; license_type=Proprietary, is_commercial=1) all carry local_path under 'C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled*.MP4' ? that directory does not exist. The folder moved to 'C:\Users\avidu\Projects\khelsutra-guru\Annotation Setup\' whose 'Golden Labelled\' now contains ONLY the six .rallies.csv label files; zero copies of the MP4s exist anywhere in the workspace (searched). The committed export artifact data/eval_export/manifest.json repeats the same 6 dead absolute paths. Code consequence at sports-data-collector/src/cli/curate.py:679-684: export of a curated video with missing local_path emits md5=None plus a warning, and src/core/manifests.py:23-26 documents md5 as "the indexer's join key" ? so any fresh export of the entire owned golden corpus loses the deterministic MD5 join to badminton-highlight-indexer. Data-loss half CONFIRMED as single-copy but the "no documented backup" wording is wrong: (a) verified live via read-only 'gcloud storage ls' (2026-07-02): gs://khelsutra-rally-corpus/videos/ has 10 objects and NONE of the six originals, while labels/ DOES hold all six (2026-06-08_adarsh_avi_singles.rallies.csv etc.) and trajectories are in the bucket ? annotation work-product is safe, source footage is not; (b) Google Drive golden folder (G:\My Drive\Sharing\Annotation Videos\Golden Label Videos Request - June 15) holds only the 7 proxy clips, not the original six; (c) the sole remaining copy is documented at F:\[Khelsutra] GoPro Hero Black 12\KheISutra\Training\Golden Labelled\ per badminton-highlight-indexer/scratch/copy_videos_to_gcs.ps1:31-36 and docs/DATA_IN_GCS.md:153 (§6 Phase-2) ? and the F: drive is NOT currently mounted (only C: and G: exist), so even that copy's integrity is unverifiable today, and no MD5s of the six were ever persisted anywhere (eval_export/manifest.json predates the md5 field from PR #51; no proxy-registry rows).
31-- notes: Corrections to the round-1 claim: (1) Size is ~37.9 GB, not 4.9 GB ? per-clip sizes annotated in scratch/copy_videos_to_gcs.ps1:31-36 (10.9+7.7+11+6.7+1.4+0.19 GB); the ~4.9-5.2 GB figure matches the three unrelated Roora pilot videos under 'Annotation Setup/Archive/.../Annotation Videos/'. (2) "No documented backup" is false ? badminton-highlight-indexer/docs/DATA_IN_GCS.md §5/§6/§8 explicitly documents the F: archive location and a Phase-2 migration plan (tracker at doc line 246: 14/15 videos still unmigrated); but "single-copy" stands: exactly one copy of the six source videos exists, on an intermittently-mounted external HDD, unverifiable right now, with no persisted hashes. (3) Downgrade from potential P0 to P1 because labels + trajectories for all six ARE in gs://khelsutra-rally-corpus (verified), so GPU-free eval/litmus and synth-train survive a drive failure; what would be permanently lost is the proprietary source footage (no re-extraction via --trajectory-source detector, no re-annotation, no re-rendering). Sharper failure scenarios: (a) F: HDD failure = permanent loss of ~38 GB of irreplaceable owned golden footage; (b) today, 'curate export' on the owned corpus emits md5=None for all six (curate.py:679-684), silently breaking the indexer MD5 join contract added in PR #51. Correctly-scoped fix (M, cross-repo): mount F:, verify + hash the six MP4s (persist MD5s into the DB/manifest immediately ? S and decision-free), generate 1080p proxies, upload via the existing scratch/copy_videos_to_gcs.ps1 or the rclone cloud path (Phase-2, already planned), then repoint/annotate local_path (or add a gs:// video_uri column) in sports_metadata.db and refresh data/eval_export/manifest.json. Partially OWNER_GATED sub-decision only: DATA_IN_GCS.md §7.4 leaves proxies-only vs full-4K migration scope OPEN and the upload is real (though tiny, cents/month) GCS spend; the hash-and-record step and the DB path fix need no owner decision.
32-
33:### R2-7 [P2/S] ? non-git
34-**Roora vendor pilot record (paid deliverables + corrected golden labels + 3 videos) is a single unversioned local copy at risk of loss**
35-- verdict: CONFIRMED

36-- reasoning: CONFIRMED for the paid deliverables + corrected golden labels; the "3 videos" component of the claim is REFUTED (Drive copies exist). Evidence: (1) Single-copy artifacts verified on disk: raw paid CVAT deliverables at 'C:/Users/avidu/Projects/khelsutra-guru/Annotation Setup/Archive/Roora Pilot Review - 6 June 2026/' ? badminton/{badminton, badminton (1), badminton (2)}/Annotations/*.xml (~2.3 MB per-frame boxes), pickleball/*.zip (3 zips, 260 KB), tablettennis/*.zip (3 zips, 213 KB) ? plus the human-corrected golden labels 'Annotation Setup/golden_labels_badminton.csv' (2.5 KB, mtime 2026-06-06) and the rescued multisport goldens 'Annotation Setup/Trajectories/roora/{pickleball,tablettennis}_rallies.csv'. (2) Not in git: sports-data-collector tracks only the AS-DELIVERED parse (data/roora_badminton_canonical.csv, verified via git ls-files + diff ? it still has rally 2 = 0:02.503?0:32.533, rally 7 end 1:51.612, no rallies 7.5/16.5), while golden_labels_badminton.csv holds the corrections (rally 2 start?0:29, rally 7 end?1:38, missed rallies 7.5@1:46 and 16.5@3:34 added, winner?out fixes) ? those corrections exist NOWHERE else (find over C:/Users/avidu/Projects + Downloads returned only this one file). (3) Not in GCS: recursive 'gcloud storage ls -r gs://khelsutra-rally-corpus/**' contains zero objects matching roora|pilot|pickle|tablettennis|sample_long|golden_labels; labels/ holds exactly the 15 canonical corpus clips (pilot video not among them); archive/ holds only a droplet archive. (4) Not on Drive: 'G:/My Drive/Sharing' has guidelines/briefs and 'Golden Annotations' (MahadevpuraSingles only) ? no CVAT exports. (5) Outside every backup convention: badminton-highlight-indexer/docs/DATA_IN_GCS.md §5 gap table (lines 117?131) covers only the 15-clip corpus and never mentions the pilot artifacts (and its line 121/142 path 'C:\?Annotation Setup\Collect\Trajectories\' is stale ? the folder moved, no Collect dir exists); the F: backup (memory checkpoint 2026-06-22) covered only the 7 new-clip labels; docs/archives/roora-vendor/README.md archives only guidelines+ADR, not pilot results; F: is currently unmounted. REFUTED part: the 3 pilot videos have byte-size-identical cloud copies at 'G:/My Drive/Sharing/Annotation Videos/Pilot Videos/' (Sample Long Video - Baadminton.MP4 = 2,882,861,900 B; Video 2 - Pickleball.mp4 = 1,435,276,306 B; Video 3 - Tablettennis.mp4 = 871,944,715 B ? all match local), and the local 'Annotation Videos/desktop.ini' references GoogleDriveFS.exe proving Drive File Stream provenance. Owner's own memory audit had flagged the precursor ("Roora multisport CSVs live ONLY in %TEMP%? one cleanup from loss"); the rescue moved them to Annotation Setup but they remain a single unversioned C:-disk copy.

37-- notes: Correction to the finding: the at-risk payload is NOT the 4.9 GB folder ? the videos are safe on the owner's Google Drive. What is genuinely one C:-disk failure (or one folder-cleanup) from loss is ~3 MB of paid, partly irreplaceable work product: the raw vendor CVAT exports for all 3 sports (box geometry the tracked canonical CSVs discard), the human frame-by-frame correction record (golden_labels_badminton.csv + rally_review.csv ? hours of review labor and the calibration evidence behind the vendor feedback), the pickleball/tablettennis golden CSVs backing the recorded TT-0.909/PB-0.308 multisport result, and the pilot review MD/PDF record. Vendor re-export from CVAT job 4075945 is not under owner control. Failure scenario: C: SSD failure or a cleanup of the non-git 'Annotation Setup' folder silently destroys the only record of the paid pilot and its corrected labels; nothing in DATA_IN_GCS.md's migration plan would ever pick it up because the pilot artifacts are absent from its §5 gap table. Fix (S, not owner-gated for the metadata): one 'gcloud storage cp -r' of 'Annotation Setup/Archive/Roora Pilot Review - 6 June 2026/' (excluding Annotation Videos/, already on Drive) + 'Annotation Setup/{golden_labels_badminton.csv,rally_review.csv,Trajectories/roora/}' to gs://khelsutra-rally-corpus/archive/roora-pilot-2026-06/ (archive/ prefix precedent already exists; ~3?10 MB, negligible cost), then add a Roora-pilot row to DATA_IN_GCS.md §5 and fix its stale 'Annotation Setup\Collect\Trajectories' path (lines 121/142). Including the 3 pilot videos in the bucket is the only owner-gated piece (DATA_IN_GCS §7 open decision 4, ~5 GB ? \$0.10/mo) and is optional given the Drive copies.

38-

39-### R2-8 [P1/S] ? badminton-highlight-indexer

40-**Upload retention sweep tool exists but nothing schedules it ? unbounded upload growth + unenforced 7-day privacy retention on any hosted deployment (#301)**

41-- verdict: CONFIRMED

42-- reasoning: Tool exists: badminton-highlight-indexer/tools/cleanup_caches.py (T11 retention sweep, --uploads-retention-days default 7.0 at ~line 407; classify() purges uploaded videos + .partial-* in the uploads zone; landed cb7394b/#299). Nothing schedules it: workspace-wide grep for cleanup_caches hits only the tool/tests/docs and an error string (backend/pipeline/detectors/base.py:162); zero references in .github/workflows/nightly.yml, deploy/cloudrun/, deploy/digitalocean/, scripts/, and zero cron/timer/sweep references in khelsutra/deploy/droplet/ (khelsutra-engine.service has no companion timer). Backend startup only sweeps Gemini remote orphans (backend/main.py:1808 _maybe_sweep_gemini_orphans), not local uploads. Issue #301 is OPEN (viewed via gh today); no scheduling commit in git log between cb7394b and master e410136. Deployment is LIVE: khelsutra/deploy/DEPLOY_ALPHA.md line 3 "Status:

LIVE (deploy live-verified on the droplet 2026-06-22)"; uploads land via backend/api/uploads.py into security.upload_dir (/srv/khelsutra/incoming/uploads per runbook S2.2).

43-- notes: Sharper failure scenario: the alpha droplet is live and accepting uploads behind Cloudflare Access; without any scheduled `python -m tools.cleanup\_caches --apply`, a tester's footage uploaded day 0 is still on disk day 30+ (unenforced 7-day privacy retention, an explicit owner policy per the tool docstring, 2026-06-21), and output/uploads grows unboundedly until require_free_bytes (backend/api/uploads.py:93) starts rejecting new uploads (service degradation on the small CPU droplet). Mitigations that keep this P1 not P0: runbook notes testers had not yet exercised the upload CUJ (OTP login + Gemini key rotation pending as of 2026-06-22), and the owner can run the sweep manually. Fix scope (S): do NOT re-implement the sweep ? ship a systemd timer + oneshot unit (or cron line) alongside khelsutra/deploy/droplet/khelsutra-engine.service that runs `python -m tools.cleanup\_caches --apply --uploads-retention-days 7 --uploads-dir <security.upload\_dir>` daily with --json output captured to journald; default-OFF/opt-in flag so local single-user installs are unaffected; document enable+verify steps in DEPLOY_ALPHA.md; one smoke test asserting the wrapper passes the right flags (tool logic already unit-tested in tests/test_cleanup_caches.py). Issue #301's acceptance criteria already spec this exactly. Partial owner gate: shipping the timer unit + docs is not gated, but *enabling* it on the live droplet is an ops action requiring the owner's SSH access, and #301 flags one open design call (whether apply should also skip uploads referenced by non-terminal jobs rows) ? recommend defaulting to the time-window-only behavior and letting the owner ratify.

44-

45-

46-## Owner-gated (flagged, no action)

47-

48-- [P2] badminton-highlight-indexer: **P0-rename remainder: person/venue-attributable video names still committed in code and baselines** ? Defect verified in current tree: backend/eval/rally_seg_eval.py:52-54 commits MULTICOURT_CLIPS = frozenset({"mahadevpura_2", "GX010128", "Badminton_BXH_2", "Boxhill_Doubles"}) (venue-attributable), and eval_baselines/heuristic_lovo_n6.json per_video keys are adarsh_avi_singles, kushagra_singles, targetest_doubles, gbaaddy, testlarge_short, mahadevpura_singles (person first names + venues). git log on rally_seg_eval.py (latest 0c9c693 #395) and on the baseline (#77) shows no rename ever landed. docs/GOLDEN_REGRESSION_FIXTURES.md:151 and docs/BACKLOG.md:61 track this exact remainder as DEFERRED per the owner's safe-scope decision (2026-06-16) with trigger "before any open-sourcing / public commit" ? a product/strategy owner gate ? and the coordinated fix requires renaming the owner-held gitignored output/human_lovo_manifest.json in lockstep or the gen0 LOVO training floor breaks. Repo is currently PRIVATE (gh repo view: Khelsutra/badminton-highlight-indexer, isPrivate=true), so no live public exposure today. Real, unfixed debt, but remediation is explicitly owner-gated on the open-sourcing decision and owner-side data coordination.

49-

50-- [P2] khelsutra (avidullu/khelsutra): **Issue #114 stale-open: CF Access guest-login debugging against a deployment that no longer exists** ? Debt is real but the fix is an owner tracker call on the owner-gated alpha/serving track. Verified: (1) gh issue view 114 ? OPEN, createdAt == updatedAt == 2026-06-23T10:25:52Z, 0 comments; body is CF-dashboard-only triage of the Access allowlist for the droplet alpha and explicitly says "Deferred by owner 2026-06-23". (2) Cited anchor exists: deploy/DEPLOY_ALPHA.md:106 is the Allow-policy/Emails-allowlist step (line 105 is the app-domain step ? off by one). (3) Deployment gone: memory session-handoff.md:19 ? droplet claude-do-rally-test/168.144.126.176 deleted 2026-06-25, "the prior 'LIVE alpha on the droplet' ... is GONE ? alpha is paused"; api.khelsutra.guru CNAME/CF-Access flagged for cleanup. (4) Not already fixed: git log -- deploy/DEPLOY_ALPHA.md shows nothing relevant after 2026-06-23 (latest 6e4dbaf PR #113 is the unrelated GPU-VM doc); no in-repo reference to #114. (5) Nothing in-repo can fix or verify it ? the Access policy lives in the Cloudflare dashboard and the origin behind app./api. no longer exists. Fix requires owner disposition (close vs retitle), squarely under the alpha/serving go-live gate.

51-

52-

53-## Critic additions (round 1 completeness pass)

54-

55-- [P1/M] critical-fix: **Golden-corpus registry rot: all 6 owned 'curated' videos point at a deleted absolute path, and the 4.9GB source videos are single-copy with no documented backup** ('sports-data-collector/data/sports_metadata.db:0')

56- - evidence: Query of the git-tracked registry: videos rows adarsh_avi_singles, kushagra_singles, targetest_doubles, gbaaddy, testlarge_short, mahadevpura_singles (all is_commercial=1, status=curated) have local_path='C:\Users\avidu\Projects\Annotation Setup\Collect\Golden Labelled\...' and video_url=NULL. `ls /c/Users/avidu/Projects/Annotation

Setup` -> 'No such file or directory'; the data now lives at 'khelsutra-guru/Annotation Setup/Golden Labelled' (folder moved, registry never updated). src/cli/curate.py:224 registers owned videos with `video_url=None, local_path=video_path` ? raw operator-supplied absolute path, no normalization or later integrity check. ShuttleSet rows are also CWD-relative with mixed separators ('./data/videos\shuttleaset_1.mp4'). 'Annotation Setup/Archive' is 4.9G of source match videos; grep for gs://[bucket]backup across sports-data-collector docs/README returns nothing, while docs/MANUAL_INGESTION_POSITIONING.md:13 declares the collector the 'single canonical system-of-record' for golden rally labels.

57- - fix: Add a `curate verify` (or extend cleanup.py) that resolves every registered local_path and flags dangling rows; re-point the 6 golden rows to the current 'khelsutra-guru/Annotation Setup/Golden Labelled' location; store paths relative to a configurable corpus root instead of absolute machine paths; and back the 4.9GB Archive up to the existing GCS bucket, recording the gs:// URI in video_url so the registry has a durable pointer.

7. user (2026-07-02T04:57:48.619Z)

530:- [P3/S] hygiene: **1.7GB of stale yt-dlp partial-download debris in data/videos and placeholder junk rows in the committed corpus registry** (`sports-data-collector/src/crawlers/badminton\_crawler.py:231`)

531:[Omitted long matching line]

532: - fix: Delete the stale *.part/*.ytdl files; extend the crawler's failure cleanup (or cleanup.py) to glob and remove yt-dlp fragment side files for the failed target; purge or clearly mark the two 'watch?v=example' placeholder rows out of the committed registry.

594:- sports-data-collector: ingest.py records a fake placeholder video URL by default (hygiene, `src/cli/ingest.py`)

8. user (2026-07-02T04:57:55.108Z)

- fix: Add a minimal .github/workflows/ci.yml to sports-obsreport: checkout, pip install -e .[dev], pytest --cov on push/PR. This closes the library half of the reporting-stack gap independently of the engine's obsreport-lane work.

- [P3/S] hygiene: **Engine .git bloated to 281MB by ~229MB of dangling raw match-video blobs; 6.5MB yolov8n.pt is permanently in pushed history and .gitignore has no guard against re-staging raw videos** (`badminton-highlight-indexer/.gitignore:49`)

- evidence: `git count-objects -vH`: 4716 loose objects, 266.33 MiB loose vs only 2.61 MiB in packs; du .git = 281M. Four loose blobs are raw MP4s (ftypisom headers): 22.4MB, 31.7MB, 84.5MB, 90.2MB ? all unreachable even via `git rev-list --all --reflog` (dangling from an aborted `git add`), never gc-pruned. Separately, yolov8n.pt (6,549,796 bytes) is reachable history on origin/master (added 5257fd5, removed d19557e 'replace AGPL ultralytics YOLO'), so every clone carries the AGPL-era weights forever. .gitignore covers *.pt (line 5) but only two narrow mp4 patterns (lines 49-50: backend/static/highlights_*.mp4, *_compiled.mp4) ? nothing stops the raw-match-video staging accident from recurring.

- fix: Run `git gc --prune=now` in the local checkout to reclaim ~260MB, and add a broad guard (e.g. `\*.mp4` with explicit `!` exceptions for any tracked fixtures, or ignore the video scratch dirs) to .gitignore. Optionally note yolov8n.pt in docs as accepted history weight; rewriting pushed history is not worth it at 6.5MB.

- [P3/S] tech-debt: **obsreport dependency URL still points at the pre-transfer avidullu org ? installs survive only via GitHub's transfer redirect** (`badminton-highlight-indexer/pyproject.toml:68`)

- evidence: pyproject.toml:68: "obsreport @ git+ssh://git@github.com/avidullu/sports-obsreport.git@v0.1.3" (echoed in requirements.txt:35). The repo actually lives at Khelsutra/sports-obsreport ? `gh repo view avidullu/sports-obsreport` resolves to nameWithOwner Khelsutra/sports-obsreport, i.e. a transfer redirect. The redirect breaks the moment any repo is re-created under the old name, and like the already-flagged trackers pin this is a mutable tag, not a commit SHA. The local clone's own remote already uses the Khelsutra URL, so the pyproject reference is the only stale spot.

- fix: Update the dependency URL in pyproject.toml (and the requirements.txt comment) to github.com/Khelsutra/sports-obsreport.git and pin the v0.1.3 commit SHA instead of the tag, matching whatever convention the trackers-pin fix adopts.

- [P3/S] hygiene: **1.7GB of stale yt-dlp partial-download debris in data/videos and placeholder junk rows in the committed corpus registry** (`sports-data-collector/src/crawlers/badminton\_crawler.py:231`)

- evidence: data/videos contains _diag.mp4.part (354MB), _t.mp4.part (516MB), _v.mp4.part (849MB) plus Frag*.part and .ytdl markers, all dated 2026-05-26 ? 1.7G total of dead partial downloads (du -ch *.part *.ytdl). The crawler's failure cleanup (badminton_crawler.py:231-246) removes only `local\_path`, never yt-dlp's .part/.part-Frag*/.ytdl side files. The tracked sports_metadata.db also carries two placeholder rows with video_url='https://youtube.com/watch?v=example' (staging_match_non_comm, shuttlestet_12) ? test fixtures committed into the canonical registry.
- fix: Delete the stale *.part/*.ytdl files; extend the crawler's failure cleanup (or cleanup.py) to glob and remove yt-dlp fragment side files for the failed target; purge or clearly mark the two 'watch?v=example' placeholder rows out of the committed registry.

Unverified backlog (dropped at the round-1 cap; lower-ranked or duplicate)

- [P2] sports-data-collector: Stale doc pointers after archive moves (cross-repo DECISIONS.md path + 6 in-code CODE_HEALTH_PLAN references)
- [P2] sports-data-collector: curate.py is a 1531-line monolith (CLI + interactive review + import/export/proxy/licensing)
- [P2] sports-data-collector: Golden-label store is a binary SQLite file tracked in git ? unmergeable across the owner's concurrent sessions
- [P1] rally-annotator (Khelsutra/rally-annotator; local checkout C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): Content script runs a 3 Hz full-DOM shadow-piercing scan on every page the user visits, even with the panel hidden
- [P1] rally-annotator (Khelsutra/rally-annotator; local checkout C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): Locked decision LD-i8 (web language persisted + seeded from browser language) is documented as shipped but not implemented ? panel always boots in English and forgets the choice on reload
- [P2] rally-annotator (Khelsutra/rally-annotator; local checkout C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): Issue #26 open and unaddressed: CSV carries zero provenance ? labels are unattributable to their source video except by filename
- [P2] rally-annotator (Khelsutra/rally-annotator; local checkout C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): E2E never reloads a page: the two headline persistence behaviors (resume numbering from storage, locale persistence) have no end-to-end proof, and entrypoints/ are outside unit coverage
- [P2] rally-annotator (Khelsutra/rally-annotator; local checkout C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): Issue #22 still fully pending: all six translation catalogs remain machine-draft ? the declared quality gate before any public launch
- [P2] rally-annotator (Khelsutra/rally-annotator; local checkout C:/Users/avidu/Projects/khelsutra-guru/rally-annotator): 'Annotation Setup' is the un-versioned annotation-data workspace (not a code duplicate) and its Roora pilot-review artifacts appear to exist nowhere else
- [P1] sports-obsreport: Collector pins obsreport as plain PyPI requirement while the name is unclaimed on PyPI (dependency-confusion window; install broken today)
- [P2] sports-obsreport: Engine (and both repos' comments) reference obsreport via the stale pre-transfer org URL avidullu/ instead of Khelsutra/
- [P2] sports-obsreport: sports-obsreport has zero CI ? tags/pushes are cut with no automated test or lint gate
- [P2] sports-obsreport: Engine-obsreport integration is never exercised in any CI ? the ~12 reporting tests only ever run on dev machines
- [P1] non-git dirs: "Annotation Setup" + "khelsutra-screencast" (workspace root C:/Users/avidu/Projects/khelsutra-guru): Roora vendor pilot record (paid deliverables + corrected golden labels + 3 videos) is a single unversioned local copy
- [P2] non-git dirs: "Annotation Setup" + "khelsutra-screencast" (workspace root C:/Users/avidu/Projects/khelsutra-guru): khelsutra-screencast: 8 ad-hoc QA scripts and doc-cited DoD evidence are unversioned; the 5 cited checkpoint screencasts are absent from the canonical dir
- [P2] non-git dirs: "Annotation Setup" + "khelsutra-screencast" (workspace root C:/Users/avidu/Projects/khelsutra-guru): DATA_IN_GCS.md \$5 gap table contradicts its own \$7: trajectories/labels marked 'not in bucket' though migrated 2026-06-22
- [P2] non-git dirs: "Annotation Setup" + "khelsutra-screencast" (workspace root C:/Users/avidu/Projects/khelsutra-guru): Stale-open #175: OneDrive golden-data sharing plan never executed and largely superseded by the GCS bucket
- [P1] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: Vendored SPA bundle in engine is 10 web/

commits stale ? engine-served UI lacks the fps display its own API just shipped

- [P1] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: Engine PR #390 (D1 main.py split) is now CONFLICTING against master and going staler
- [P1] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: Engine CLAUDE.md 'Current state' is ~3 weeks false ? claims platform/owned-model work 'has not started'
- [P1] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: D0 droplet retired 2026-06-25 but runbooks in BOTH repos still say the deploy is LIVE on it
- [P2] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: sports-obsreport has zero CI ? no .github directory at all despite shipping versioned releases
- [P2] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: Lint/type gates exist only in the engine ? sports-data-collector CI is pytest-only, plus GH-Actions version drift across repos
- [P2] khelsutra-guru workspace (cross-repo): badminton-highlight-indexer, khelsutra, sports-data-collector, rally-annotator, sports-obsreport: sports-data-collector working tree is dirty: uncommitted .gitignore WIP + a full stale nested clone of itself
- [P0] badminton-highlight-indexer (Khelsutra): #340/#330: /api/process still re-bills Gemini on re-submit of an already-processed video ? no idempotency short-circuit (live P0 cost bug)
- [P1] badminton-highlight-indexer (Khelsutra): PR #390 (D1 main.py split, step 1) is CONFLICTING and stale ? master diverged 9 commits on main.py since it was cut
- [P1] badminton-highlight-indexer (Khelsutra): #332: NSFW pre-flight gate still absent ? non-owner alpha users can route arbitrary video to Gemini (ToS/abuse + paid-garbage risk)
- [P2] badminton-highlight-indexer (Khelsutra): #327: path-jail still blocks gs://-style ingest in hosted mode ? validate_input_path runs before scheme resolution
- [P2] badminton-highlight-indexer (Khelsutra): #388 premise is half-stale: the three job-worker functions already live in backend/api/job_worker.py (since PR #234, 06-18) ? remaining scope is shim removal + test patch-target migration
- [P2] badminton-highlight-indexer (Khelsutra): #384: coverage floor 70 and offline guard lanes still duplicated verbatim between ci.yml and nightly.yml
- [P2] badminton-highlight-indexer (Khelsutra): #301: upload retention sweep tool exists but nothing schedules it ? unbounded upload growth on any hosted deployment
- [P2] badminton-highlight-indexer (Khelsutra): #349 (+#331 overlap): WASB streaming decode still forced single-threaded ? L4 GPU ~80% idle; #331 additionally needs a refresh (droplet baseline retired, its lever 3 already shipped)
- [P2] badminton-highlight-indexer (Khelsutra): #328: per-video runlog/trace endpoint still missing ? obsreport exists server-side but is not user-accessible
- [P1] Khelsutra/badminton-highlight-indexer: No timeout on any serving-path ffmpeg subprocess ? one hung ffmpeg permanently wedges the single-worker job queue
- [P1] Khelsutra/badminton-highlight-indexer: MD5 segmentation reuse still unimplemented ? every re-submit of a byte-identical video re-bills Gemini and re-pays the full CPU pipeline
- [P1] Khelsutra/badminton-highlight-indexer: WASB streaming decode still single-threaded (num_workers forced to 0), no NVDEC anywhere ? L4 GPU ~80% idle during detection
- [P1] Khelsutra/badminton-highlight-indexer: Upload retention sweep still manual ? 7-day privacy retention window unenforced without an operator (#301 open)
- [P2] Khelsutra/badminton-highlight-indexer: Remote-input scratch copies (rally-input-*/rally-worker-*) are never reclaimed ? by-design temp leak outside all cleanup tooling

Unverified P3 pool

- badminton-highlight-indexer: regression.py silently passes when no baseline exists (audit B8, still open) (critical-fix, `backend/eval/regression.py`)
- badminton-highlight-indexer: Chunked upload: concurrent/proxy-retried PUT of the same part index passes the sequence check twice and interleave-appends ? corrupt assembly detected only by part-count mismatch, forcing a full multi-GB re-upload (test-gap, `backend/api/uploads.py`)
- badminton-highlight-indexer: A failed mark_job_done/mark_job_failed write leaves a job in 'running' forever (writers swallow to False, worker ignores the return) ? pollers hang until the next restart's stale sweep (test-gap, `backend/api/job\_worker.py`)
- badminton-highlight-indexer: 4 stale remote branches from merged/superseded PRs never deleted (hygiene, `.git`)
- badminton-highlight-indexer: D13 is DONE but yolo_hybrid.py docstring still claims tracking

uses supervision's MIT ByteTrack (licensing-accounting text now wrong) (doc-rot, `backend/pipeline/segmenters/yolo\_hybrid.py`)

- badminton-highlight-indexer: ablation PDF exporter (backend/eval/ablation_pdf.py) has zero CI coverage ? its 2 tests always skip for missing reportlab/pypdf (test-gap, `tests/test\_ablation.py`)
- badminton-highlight-indexer: CI test output hides skips and slow tests: -q with no -ra / --durations anywhere (hygiene, `run\_all\_tests.py`)
- badminton-highlight-indexer: ruff target-version pinned to py38 while the package requires Python >=3.10 and CI runs 3.13 (hygiene, `ruff.toml`)
- badminton-highlight-indexer: JWKS fetch does not enforce https on cf_team_domain ? a http:// misconfig would fetch signing keys over cleartext (security, `backend/api/auth.py`)
- badminton-highlight-indexer: GPU-box runbook instructs pasting the raw Gemini key inline in a shell command (lands in root bash history) (hygiene, `deploy/gcp/README.md`)
- badminton-highlight-indexer: eval_baselines/reelcount_manifest.json hardcodes gs://khealsutra-rally-corpus URIs outside the buckets.env single-source (tech-debt, `eval\_baselines/reelcount\_manifest.json`)
- khealsutra (avidullu/khealsutra) ? SPA web/ + marketing site/: Open UX issues #99 and #103 look overtaken by the 2026-07-01 PR wave but were never reconciled (hygiene, `khealsutra/web/src/components/VideoLibrary.tsx`)
- khealsutra (avidullu/khealsutra) ? SPA web/ + marketing site/: web/ SPA has no linter (no ESLint config or lint script) ? react-hooks correctness classes are unguarded (hygiene, `khealsutra/web/package.json`)
- khealsutra (avidullu/khealsutra): Issue #102 fully open: no &debug= URL-param handling exists in the SPA (test-gap, `web/src/App.tsx`)
- khealsutra (avidullu/khealsutra): Issue #72 fully open: Telugu locale not started (hygiene, `web/src/i18n/index.ts`)
- sports-data-collector: ingest.py records a fake placeholder video URL by default (hygiene, `src/cli/ingest.py`)
- rally-annotator (Khealsutra/rally-annotator; local checkout C:/Users/avidu/Projects/khealsutra-guru/rally-annotator): Local git remote still points at pre-transfer avidullu/rally-annotator instead of canonical Khealsutra/rally-annotator (hygiene, `C:/Users/avidu/Projects/khealsutra-guru/rally-annotator/.git/config`)
- sports-obsreport: README status section is three phases and three releases stale (doc-rot, `C:/Users/avidu/Projects/khealsutra-guru/sports-obsreport/README.md`)
- sports-obsreport: _plural() mishandles common nouns (e.g. item_noun="match" renders "3 matchs") in customer-facing summaries (critical-fix, `C:/Users/avidu/Projects/khealsutra-guru/sports-obsreport/src/obsreport/render/markdown.py`)
- non-git dirs: "Annotation Setup" + "khealsutra-screencast" (workspace root C:/Users/avidu/Projects/khealsutra-guru): 0.5 GB duplicate video object in the corpus bucket awaiting the prescribed verify-dupe-then-delete (hygiene, `badminton-highlight-indexer/docs/DATA\_IN\_GCS.md`)
- khealsutra-guru workspace (cross-repo): badminton-highlight-indexer, khealsutra, sports-data-collector, rally-annotator, sports-obsreport: 13 remote + 7 local branches across the workspace are merged/closed debris (hygiene, `khealsutra`)
- khealsutra-guru workspace (cross-repo): badminton-highlight-indexer, khealsutra, sports-data-collector, rally-annotator, sports-obsreport: SPA comments hard-code engine main.py line numbers that are already wrong and about to get worse (doc-rot, `khealsutra/web/src/api/types.ts`)

9. user (2026-07-02T04:58:46.986Z)

349 C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/pr55.diff

10. user (2026-07-02T04:58:55.785Z)

```

README.md
cleanup.py
config
data
docs
pyproject.toml
src
tests
aef4621 feat(curate): add `verify` local_path audit; fix golden-corpus registry rot
adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)

```

6f67d65 feat(import): callable import_annotations + a maturity gate (P9) (#52)
c281f82 feat(export): add source file MD5 to the eval/training manifest (P9) (#51)
a152cf5 fix(cvat): warn on rows silently skipped by load_canonical_csv (#50)

11. assistant (2026-07-02T04:58:56.627Z)

You've hit your session limit · resets 1:10pm (Asia/Calcutta)