

claude-windows / c1fcb4aa-537d-41fd-858e-53f93349bf5a

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\subagents\workflows\wf_7049e1ed-b93\agent-a6180dc543e3f9628.jsonl`  
- SHA-256: `97ad3bbfb62908198605e066d6a27ce72c8a584b8e482ca2bfef52fb9d01bc67`  
- Source modified: `2026-07-02T04:54:07+00:00`  
- Imported at: `2026-07-05T16:48:24+00:00`  
- project: `wf_7049e1ed-b93`  
- session_id: `c1fcb4aa-537d-41fd-858e-53f93349bf5a`
```

Transcript

1. user (2026-07-02T04:51:09.996Z)

You are an adversarial code reviewer. Today is 2026-07-02. Review PR #54 in the repo at C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector (gh pr view 54, gh pr diff 54; the working copy is in worktree C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-gitignore-artifacts ? read code there for context).
The PR claims to fix: R1-39 (uncommitted .gitignore hunk) (evidence in C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/AUDIT_REPORT_khelsutra-guru_2026-07-02.md ? Grep for the IDs; honor the verifier notes' scoping corrections).
Hunt REAL defects: correctness bugs in the diff, broken/vacuous tests, behavior changes beyond the brief, staging accidents (files that should not be in the diff ? check the file list!), convention violations (default-ON when brief says OFF, missing schema-version check, CWD assumptions), and whether the fix actually addresses the finding. READ-ONLY ? do not modify or comment on GitHub. blocker = must not merge; major = fix before merge; minor = nice-to-have.

2. user (2026-07-02T04:51:18.099Z)

```
title: docs(gitignore): record that model-artifact bundles never live in this repo (ADR-008)  
state: OPEN  
author: avidullu (Avi)  
labels:  
assignees:  
reviewers:  
projects:  
milestone:  
number: 54  
url: https://github.com/Khelsutra/sports-data-collector/pull/54  
additions: 14  
deletions: 0  
auto-merge: disabled  
--
```

What / why

Resolves audit finding ****R1-39**** (2026-07-02 khelsutra-guru audit): the main checkout of this repo has carried an ****uncommitted**** `.gitignore` hunk for ~3 weeks (mtime 2026-06-10) that appended a blanket `artifacts/` ignore with a comment citing ADR-008 and declaring "the reviewable surface is each bundle's manifest.json".

That hunk was wrong on two counts:

- **Wrong repo.**** ADR-008 "Owned-model topology" does not exist in this repo (`docs/DECISIONS.md` has only ADR-001/ADR-002); it lives in the indexer (`badminton-highlight-indexer/docs/archives/decisions/DECISIONS.md`). Its actual decision ? confirmed against the

indexer's `docs/CODE\_MAP.md` \$0 placement table ? puts artifact bundles at
`artifacts/<name>/<ver>/` **in the indexer repo** ("manifest committed (provenance), weights
gitignored"), and pre-commits the Gen-2 split to a **new** repo (working name `sports-temporal-
models`) ? never the collector. This repo's role in ADR-008 is the data feed: `curate export` ?
`data/eval\_export\*/`, already ignored.
2. **Self-contradictory.** A blanket `artifacts/` ignore hides the very `manifest.json` the
comment declared reviewable.

Since bundles never live here, this PR adds **no ignore pattern**. It documents the decision in
`.gitignore` itself (matching the file's existing commented-section style) so the question
doesn't resurface, and records the correct un-ignore recipe (`artifacts/\*\*` + `!artifacts/\*\*/` +
`!artifacts/\*\*/manifest.json`) should bundles ever land here.

Note for the orchestrator: the main checkout (`C:/Users/avidu/Projects/khelsutra-
guru/sports-data-collector`) still carries the stale WIP `.gitignore` hunk. This PR was authored
fresh from `origin/master` in an isolated worktree; the stale hunk is superseded by this PR and
should be **discarded** (`git checkout -- .gitignore`) after merge.

Verification

- Comment-only change: `git diff` is 14 added lines, all `#`-prefixed or blank ? no ignore-rule
change, so no tests to add (no behavior changed).
- `git check-ignore -v` in the worktree: existing patterns still match (`data/videos/x.mp4`,
`data/temp/t.csv`, `data/\*.backup-\*.db`, `\*.report.json`) and `artifacts/gen0/1.0/manifest.json`
/ `weights.pt` are **not** ignored (exit 1), i.e. the file parses and semantics are identical to
`origin/master`.
- ruff/mypy: not applicable (no Python touched).
- CI (pytest) covers the unchanged code paths.

? Generated with [Claude Code](https://claude.com/claude-code)

=====DIFF=====

diff --git a/.gitignore b/.gitignore

index 52dac92..09d054b 100644

--- a/.gitignore

+++ b/.gitignore

@@ -53,3 +53,17 @@ data/temp/

Local DB backups (e.g. data/sports_metadata.backup-YYYY-MM-DD.db) ? never commit
data/*.backup-*.db

+
+## ---- Model artifact bundles: NOT hosted in this repo (decision, 2026-07-02) ----

+## ADR-008 "Owned-model topology" (badminton-highlight-indexer/docs/archives/
+## decisions/DECISIONS.md) places artifact bundles in the INDEXER repo at
+## artifacts/<name>/<ver>/ ? weights gitignored (private Releases), manifest.json
+## committed for provenance ? and pre-commits the Gen-2 split to a NEW repo
+## (working name sports-temporal-models), never this one. This repo is the data
+## feed (`curate export` ? data/eval_export*/, ignored above), so there is
+## intentionally no `artifacts/` pattern here. If bundles ever DO land in this
+## repo, ignore weights while keeping manifests reviewable, e.g.:

+## artifacts/**

+## !artifacts/**/

+## !artifacts/**/manifest.json

+## ? never a blanket `artifacts/`, which would hide the manifests too.

3. user (2026-07-02T04:51:18.227Z)

384-

385-### R1-38 [P2/M] test-gap ? khelsutra (avidullu/khelsutra) ? SPA web/ + marketing site/

386-**SPA e2e CUJ never runs in CI and the mock engine is hand-synced ? the mock-vs-real drift
bug-class already shipped once**

387-- anchor: `khelsutra/.github/workflows/test.yml:31` · tracked in: e2e/mock-engine.mjs:5
('later drive CI CUJ replays') · finder: khelsutra:code

388-[Omitted long context line]

389-- why it matters:

```

390-- proposed fix: Add a fourth CI job (mirroring the existing layout job's Playwright setup)
that builds web/ and runs the e2e CUJ specs against mock-engine.mjs; separately add a thin
contract check pinning the mock's response shapes to web/src/api/types.ts (or an engine-side
snapshot of _video_summary/capabilities) so mock-vs-real drift fails a build instead of
shipping.
391-[Omitted long context line]
392-[Omitted long context line]
393-
394:### R1-39 [P2/S] hygiene ? sports-data-collector
395-**Uncommitted .gitignore hunk (artifacts/) stale ~3 weeks and self-contradictory vs the ADR
it cites**
396-- anchor: `.gitignore:59` · tracked in: ? · finder: sdc:audit
397-[Omitted long context line]
398-- why it matters:
399-- proposed fix: Either commit the hunk as a tiny PR with a pattern that actually matches the
stated intent (e.g. 'artifacts/**' + '!artifacts/', '!artifacts/**/',
'!artifacts/**/manifest.json'), or discard it if model artifact bundles will only ever live in
the indexer repo (per ADR-008 / indexer CODE_MAP, artifacts/<name>/<ver>/ is an indexer-repo
location). Owner should confirm which repo hosts bundles before committing.
400-[Omitted long context line]
401-[Omitted long context line]
402-
403:### R1-40 [P2/M] ci-infra ? sports-data-collector
404-**CI is pytest-only: no ruff/mypy/coverage gate, and default ruff already fails with 24
errors**

```

4. user (2026-07-02T04:51:35.253Z)

On branch docs/gitignore-artifact-bundles-decision
Your branch is up to date with 'origin/docs/gitignore-artifact-bundles-decision'.

```

nothing to commit, working tree clean
ab6f68b docs(gitignore): record that model-artifact bundles never live in this repo (ADR-008)
adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)
6f67d65 feat(import): callable import_annotations + a maturity gate (P9) (#52)
c281f82 feat(export): add source file MD5 to the eval/training manifest (P9) (#51)
a152cf5 fix(cvat): warn on rows silently skipped by load_canonical_csv (#50)
docs/gitignore-artifact-bundles-decision
=====
ab6f68b docs(gitignore): record that model-artifact bundles never live in this repo (ADR-008)
=====DIFF STAT=====
 .gitignore | 14 ++++++++
 1 file changed, 14 insertions(+)

```

5. user (2026-07-02T04:51:35.566Z)

```

M .gitignore
=====STASH/WIP=====
warning: in the working copy of '.gitignore', LF will be replaced by CRLF the next time Git
touches it
diff --git a/.gitignore b/.gitignore
index 52dac92..ac1420f 100644
--- a/.gitignore
+++ b/.gitignore
@@ -53,3 +53,7 @@ data/temp/

# Local DB backups (e.g. data/sports_metadata.backup-YYYY-MM-DD.db) ? never commit
data/*.backup-*.db
+
+# Model artifact bundles: weights are NEVER committed (private Releases); the
+# reviewable surface is each bundle's manifest.json (ADR-008).
+artifacts/

```

6. user (2026-07-02T04:51:45.202Z)

```
1 # ---- Python ----
2 __pycache__/
3 *.py[cod]
4 *$py.class
5 *.egg-info/
6 .eggs/
7 build/
8 dist/
9 .installed.cfg
10 *.egg
11
12 # ---- Test / coverage ----
13 .pytest_cache/
14 .coverage
15 .coverage.*
16 htmlcov/
17 .tox/
18 .cache/
19
20 # ---- Virtual environments ----
21 .venv/
22 venv/
23 env/
24 ENV/
25
26 # ---- Secrets / local config ----
27 .env
28 .env.*
29
30 # ---- OS / editor cruft ----
31 .DS_Store
32 Thumbs.db
33 .idea/
34 .vscode/
35 *.swp
36
37 # ---- Project: large media kept outside git (see README) ----
38 # The video corpus is hundreds of MB+; only metadata (the SQLite DB) and the
39 # small source annotation CSVs are version-controlled. Re-fetch videos from the
40 # stored video_url / local_path per the manifest.
41 data/videos/
42
43 # ---- Generated eval/training exports (regenerated via `curate export`) ----
44 data/eval_export*/
45
46 # ---- Generated temp artifacts (crawler scratch space; regenerated each run) ----
47 data/temp/
48
49 # ---- Per-video observability reports (written next to videos / inputs at runtime) ----
50 *.report.json
51 *.report.md
52 *.summary.md
53
54 # Local DB backups (e.g. data/sports_metadata.backup-YYYY-MM-DD.db) ? never commit
55 data/*.backup-*.db
56
57 # ---- Model artifact bundles: NOT hosted in this repo (decision, 2026-07-02) ----
58 # ADR-008 "Owned-model topology" (badminton-highlight-indexer/docs/archives/
59 # decisions/DECISIONS.md) places artifact bundles in the INDEXER repo at
60 # artifacts/<name>/<ver>/ ? weights gitignored (private Releases), manifest.json
61 # committed for provenance ? and pre-commits the Gen-2 split to a NEW repo
62 # (working name sports-temporal-models), never this one. This repo is the data
63 # feed (`curate export` ? data/eval_export*/, ignored above), so there is
64 # intentionally no `artifacts/` pattern here. If bundles ever DO land in this
65 # repo, ignore weights while keeping manifests reviewable, e.g.:
```

```

66      # artifacts/**
67      # !artifacts/**/
68      # !artifacts/**/manifest.json
69      # ? never a blanket `artifacts/`, which would hide the manifests too.
70

```

7. user (2026-07-02T04:51:46.934Z)

```

C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\PIPELINE_ROLE.md:13:-
`badminton-highlight-indexer/docs/DECISIONS.md` ? ADR-002 (this bridge), ADR-003
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\PIPELINE_ROLE.md:14:
(acquisition policy), ADR-004 (offline learned segmenter).
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\PIPELINE_ROLE.md:33:
*segmentation*, not the frame-level shuttle detector (see the indexer's ADR-002).
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\PIPELINE_ROLE.md:38: hard-won
pitfalls (SABR?144p, Chrome DPAPI, rate-limiting) are in the indexer's ADR-003.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\INGESTION_PIPELINE.md:35:#
(ADR-002 ? downscaled + scrub-friendly so the vendor can work at frame step 1;
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\INGESTION_PIPELINE.md:167:the
original. Full rationale: ADR-002 in `docs/DECISIONS.md`.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\DECISIONS.md:8:## ADR-001 ?
Keep explicit per-sport code instead of a sport registry (for now)
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\DECISIONS.md:41:- If the cost
rises, ADR-001 should be revisited and superseded.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\DECISIONS.md:65:## Checklist:
adding a new sport today (under ADR-001)
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\DECISIONS.md:76:## ADR-002 ?
Annotation proxies are generated in-repo (`curate prep-annotation-proxy`), not by a one-off tool
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:6:> **Relation to existing docs:** peer-of
`docs/DECISIONS.md` (ADR-001, ADR-002); the de-dup work items (T1/A4) are explicitly **gated
behind an ADR-001 revisit** (see D2).
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:8:> **Honesty note:** every P0/P1 claim below is
tagged `[verified]` (checked against current `master` source) or `[researched]` (needs sign-
off). ADR-001 is `[verified]`; the refactor items that conflict with it are marked ? gated.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:14:[Omitted long matching line]
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:22:**Read first:** `README.md`, `docs/DECISIONS.md`
(ADR-001 + ADR-002), `docs/PIPELINE_ROLE.md`, `src/core/db.py`, `src/core/schemas/base.py`,
`src/parsers/cvat/rally_cvat.py`.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:29:| D2 | **Honor ADR-001.** Do NOT plan the per-
sport de-dup refactor (T1, A4) for execution. List it as owner-gated; only revisit if an ADR-001
trigger fires. | ADR-001 `[verified]` | T1/A4 are tracked but ?; the "Future design sketch" in
ADR-001 is the starting point if revisited. |
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:35:[Omitted long matching line]
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:39:- **`[verified]`** ADR-001 exists and explicitly
keeps per-sport code; its "revisit triggers" are: sport count > ~8, a 4th+ rally-identical
sport, or a cross-cutting feature requiring edits to every per-sport branch. **None have
fired.** ? T1/A4 stay gated.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:48:No new architecture is introduced by the P0/P1
work ? each item is a localised fix. The only structural proposal is **A4 (sport registry)**,
which is the same idea ADR-001 already sketches and defers; it is **not planned here**.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:68:- T1/A4 (de-dup) are **not executed** ? the
duplication remains until ADR-001 is revisited.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:132:| T1 / A4 | Per-sport de-dup refactor / sport
registry. | **Conflicts with ADR-001** (deliberately keeps explicit per-sport code). T2 already
removed the worst schema-level duplication. Revisit only if an ADR-001 trigger fires. |

```

```

C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:151:- [x] ADR-001 is left untouched (T1/A4 dropped
per D8) ? never silently contradicted.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:167:- `docs/DECISIONS.md` ? ADR-001 (per-sport code,
? to T1/A4), ADR-002 (proxy contract).
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:175:- `2026-06-24` ? Created from the 2026-06-24
full code-health audit. Docs-only PR (D6); no code items closed. Tier S?D sequencing + ADR-001
gating recorded as decisions D1?D6.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:180:- `2026-06-24` ? **Tier D triaged (D8).** T1/A4
dropped (ADR-001 stands); A2/A5/A6 promoted to active; A1/A3 parked in `docs/NEXT_STEPS.md`. All
OQs resolved. Final 3 items (A2/A5/A6) remain before DONE.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\archives\CODE_HEALTH_PLAN.md:181:- `2026-06-24` ? **PLAN DONE.** A6 merged (#46);
A2 merged (#47); A5 merged (#48). All active Tier D items complete. A1/A3 remain parked in
NEXT_STEPS with triggers; T1/A4 dropped per ADR-001. Final: 280 tests / 82% coverage (was 193 /
81%). Status ? DONE; archived to `docs/archives/`.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\annotation\ROORA_RUNBOOK.md:8:ADR-002 in [ `../DECISIONS.md` ](../DECISIONS.md).
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\docs\annotation\ROORA_RUNBOOK.md:197:grid. Full decision record: ADR-002 in
[ `../DECISIONS.md` ](../DECISIONS.md);

```

8. user (2026-07-02T04:51:54.137Z)

```

Found 19 files
docs\CODE_MAP.md
docs\NEXT_STEPS.md
training\gen0\harness.py
tests\test_gt_loader.py
backend\eval\rally_seq_proto.py
backend\features\rally_seq.py
backend\eval\calibrate_wasb.py
backend\eval\gt_loader.py
backend\eval\annotations.py
docs\archives\past_projects\DECOUPLED_COMPUTE\04-BUILD-AND-TEST-PLAN.md
docs\archives\past_projects\DECOUPLED_COMPUTE\02-COMPUTE-MAP-AND-CONTRACT.md
docs\archives\past_projects\DECOUPLED_COMPUTE\01-EXISTING-PLANS-AND-GAP.md
docs\archives\decisions\DECISIONS.md
docs\COMPUTE_DECOUPLED_SERVING\README.md
training\__init__.py
tests\test_import_smoke.py
backend\features\__init__.py
training\README.md
training\requirements-train.txt

```

9. user (2026-07-02T04:51:59.463Z)

```

385-efficient, and *each one generating the very golden labels an end-to-end model will need*.
When the corpus is
386-healthy, we train the unified model. **Option A is not a detour from Option B; it is how we
earn it.**
387-
388:## ADR-008: Owned-model topology ? artifact boundary now, repo split at Gen-2
389-
390-**Date:** 2026-06-10
391-**Status:** Accepted (owner-ratified after an adversarial design panel)
--
495-- Local-first (owner-requested): removes cloud variables while the new job/upload surface
496- stabilizes; deployment becomes a configuration exercise against a proven build.
497-- Planning docs stay canonical in the engine repo for now (single decision ledger; the copy-
drift
498: failure mode is documented in ADR-008's rationale) with pointer READMEs in `khelsutra`.

```

```

499-
500-### Consequences
501-
--
688-Date: 2026-06-21 . Status: ? RATIFIED (owner, 2026-06-21 ? M0 design sign-off;
build M1+). Refines ADR-011 + ADR-012 (it realizes their deferred M3/M4; it does not
supersede them); builds on ADR-013 (GCP sole compute stack). Full design:
`docs/COMPUTE_DECOUPLED_SERVING/`.
689-
690-### Evolution / lineage (trace the idea)
691:[Omitted long matching line]
692-
693-### Decision
694-

```

10. user (2026-07-02T04:52:05.364Z)

```

388    ## ADR-008: Owned-model topology ? artifact boundary now, repo split at Gen-2
389
390    Date: 2026-06-10
391    Status: Accepted (owner-ratified after an adversarial design panel)
392
393    ### Context
394
395    M0.4 (Gen-0 TAL train?eval?ship-gate harness, productionizing
`backend/eval/rally_seq_proto.py`
396    over ActionFormer/ASFormer) and M0.1 (corpus spine + event-index contract) are greenlit.
The
397    question: should the model/training work live in this repo, or in a separate
repo whose
398    output the indexer consumes as another segmenter (like WASB and Gemini)? The owner's
instinct
399    was isolation ("iterate on model and processing independently"), citing the collector
split as
400    precedent.
401
402    ### Decision
403
404    Draw the boundary at the ARTIFACT CONTRACT now; split the repo at Gen-2, as a pre-
committed,
405    trigger-based event ? not today.
406
407    1. Training lives in this repo as a top-level `training/` tree (sibling of
`backend/`),
408    isolated executably, not socially: its own `training/requirements-train.txt`
(heavy/
409    training-only deps never enter `requirements.txt`), and the import-smoke suite
asserts
410    `import backend.main` loads no `training.` module.
411    2. The featurizer is single-source-of-truth: promoted out of `rally_seq_proto` into
412    `backend/features/rally_seq.py` (pure numpy/scipy, versioned schema constant),
imported by
413    BOTH the training harness and the future TAL serving segmenter. Train == serve by
414    construction ? the fps-normalization bug (fixed 2026-06-10) is the skew class this
kills.
415    3. Serving consumes only a versioned ARTIFACT BUNDLE (weights gitignored / private
GitHub
416    Releases + a manifest carrying: feature-schema version, normalization, calibration/
417    post-processing, training-corpus ref + GT hash, consent + license attestation incl.
the
418    WASB C3 checkpoint status). The loader/gate hard-fails on schema mismatch ? the
419    `regression.py` incommensurability-guard pattern.
420    4. The ship-gate is owned by the product side (extends `backend/eval/regression.py`
+
421    `metrics.match_intervals`): the trainer runs it as pre-flight but never grades

```

```

itself.
422     The gate refuses commercial ship while C3 is unresolved, regardless of F1.
423 5. **Pre-committed split trigger** ? extract `training/` into a separate repo (working
name
424     `sports-temporal-models`, to be finalized with the pending repo rename) **when**:
425     (a) Gen-2 arrives (ms-swift / Qwen3-VL / cloud-GPU training ? an alien dependency
tree), OR
426     (b) a second consumer of the trained artifacts appears (e.g. the khelacharya layer),
427     AND (c) the harness's only data input is the collector's M0.1 export. By then the
artifact
428     contract is proven and extraction is a file move across a known seam.
429
430 **Prerequisite frozen first: the golden-label/corpus contract.** One canonical GT loader
(the
431     repo carried two incompatible ones: `annotations.load_annotations` `start,end` vs
432     `calibrate_wasb.load_golden_gts` `start_time,end_time`, the latter feeding 7 modules) +
a
433     round-trip test against the collector's `curate export` format. Every ship-gate quantity
434     (gt_hash, LOVO numbers, the future consent fingerprint) flows through this boundary.
435
436 ### Rationale (the decisive, code-verified facts)
437
438 - **The "just another runner like WASB/Gemini" analogy fails at the seam location.**
WASB's
439     contract sits at the video boundary (video ? trajectory CSV); Gemini's at video ?
segments.
440     The Gen-0 head consumes **features this repo computes** from the trajectory ? the seam
is
441     mid-pipeline. Since the indexer is an app (not pip-installable), a separate repo today
could
442     share that code only by **copying** it ? this project's one empirically observed
failure
443     mode (the GT-header fork; the duplicated TUNED constants).
444 - **The in-house separate-repo precedent is unproven for load-bearing deps**: the
`obsreport`
445     pin has never been live (commented in requirements.txt, repo unpublished, CI skips via
446     importorskip). It works because its absence is a harmless no-op ? a property a
featurizer
447     can never have.
448 - **Gen-0 has no environmental forcing function** (trains in minutes, ~4.5 GB, MIT pip-
able
449     deps, on the owner's WSL/2070S box); Gen-2 genuinely does. Splitting then is a
scheduled
450     event, not an overruled instinct.
451 - The prerequisite work for a clean split (featurizer promotion, loader consolidation,
452     manifest) is **identical** to the in-repo design ? doing it in place first makes the
453     eventual extraction cheap and safe. This is "envision now, build later" applied to
topology:
454     **the split is driven by productionization needs, not code-isolation aesthetics**
(owner's
455     framing, ratified).
456
457 ### Consequences
458
459 - `backend/` may never import `training.` (CI-enforced). Training tests use **synthetic
460     fixtures only** until the M0.1 consent ledger exists (no real trajectories/labels
committed).
461 - Weights are never committed; artifact manifests are the reviewable/committable
surface.
462 - The collector remains system-of-record for labels + consent; the indexer owns scoring
and
463     the ship/no-ship decision; manifests pin both by hash.
464 - Revisit triggers are explicit (above); if agent-driven development repeatedly breaches
the
465     import wall despite CI, that is evidence the social+test boundary failed and the split

```


466 should happen early.
467

11. user (2026-07-02T04:52:16.570Z)

```
32-| A test | `tests/test_*.py` | pytest auto-discovers; the full-suite lane has a **70%  
coverage floor**. Pure/offline + mocked (no GPU/WSL/network). |  
33-| A doc (design, plan, living log) | `docs/` | Index it in `docs/README.md`. Only **terminal-  
state** (DONE/REJECTED) work moves to `docs/archives/`. |  
34-| A throwaway repro / debug script | `scratch/` | **Gitignored** + excluded from ruff/pytest.  
Never imported by app code. |  
35-| A model artifact (weights + manifest) | `artifacts/<name>/<ver>/` | ADR-008: **manifest  
committed** (provenance), **weights gitignored**. |  
36-| Eval baseline / leaderboard JSON | `eval_baselines/` | Tracked (survives a clean checkout);  
the no-regression gate reads it. |  
37-| **Committed regression fixture** (video-free golden) | `eval_baselines/fixtures/<slug>/` |  
Tracked, **LF-pinned** (`.gitattributes`). Synthetic Tier-0 (or owner-cleared, de-attributed  
real). Drives the `golden_fixtures` characterization gate with **no video/GPU/DB**. See  
`docs/GOLDEN_REGRESSION_FIXTURES.md`; mint via `python -m backend.eval.golden_fixtures --freeze-  
synthetic` / `--freeze-real` (refusal-gated). |  
38-| Pipeline run outputs (reels, trajectories, DBs, frames) | `output/` | **Gitignored.** Per-  
video isolation is the tracked **`PER_VIDEO_WORKSPACE.md`** project (P0 helper shipped #264;  
supersedes the archived `PER_VIDEO_OUTPUT_LAYOUT`). |
```

12. user (2026-07-02T04:52:45.806Z)

```
0  
non-comment added lines above (0 = inert)  
====check-ignore semantics (current PR state)====  
.gitignore:41:data/videos/          data/videos/x.mp4  
.gitignore:47:data/temp/           data/temp/t.csv  
.gitignore:55:data/*.backup-*.db    data/sports_metadata.backup-2026-01-01.db  
.gitignore:50:*.report.json        foo.report.json  
NOT IGNORED: artifacts/gen0/1.0/manifest.json  
NOT IGNORED: artifacts/gen0/1.0/weights.pt
```

13. user (2026-07-02T04:52:50.181Z)

```
add '.gitignore'  
add 'artifacts/gen0/1.0/manifest.json'  
!! artifacts/gen0/1.0/weights.pt  
!! artifacts/top-manifest.json  
?? .gitignore  
?? artifacts/  
====untracked files git would add====  
.gitignore  
artifacts/gen0/1.0/manifest.json
```

14. user (2026-07-02T04:52:50.944Z)

```
C:\Users\avidu\Projects\khehsutra-guru\sports-data-collector\src\reporting\__init__.py:38:  
<id>_report.* artifacts). Used by validate-delivery / convert-cvat. Never raises.""")  
C:\Users\avidu\Projects\khehsutra-guru\sports-data-collector\src\cli\curate.py:1111:      #  
colocated with the existing artifacts. Never raises.  
C:\Users\avidu\Projects\khehsutra-guru\sports-data-  
collector\src\parsers\cvat\rally_cvat.py:590:# These turn a (records, issues) pair into the  
artifacts a human reviews:      #
```

15. user (2026-07-02T04:53:32.114Z)

```
no local venv  
..... [ 49%]  
..... [ 74%]  
..... [ 99%]
```

[100%]

289 passed, 2 skipped in 7.79s

16. user (2026-07-02T04:54:07.865Z)

Structured output provided successfully