

claude-windows / c1fcb4aa-537d-41fd-858e-53f93349bf5a

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\subagents\workflows\wf_41c8dc8a-8a4\agent-a222162fd72861c99.jsonl`  
- SHA-256: `7542c59db01651f9a678ed3bd84da44fccd1e7d35f297e94b449dd381387b1d8`  
- Source modified: `2026-07-02T02:58:05+00:00`  
- Imported at: `2026-07-05T16:48:24+00:00`  
- project: `wf_41c8dc8a-8a4`  
- session_id: `c1fcb4aa-537d-41fd-858e-53f93349bf5a`
```

Transcript

1. user (2026-07-02T02:47:11.983Z)

CONTEXT (read carefully):

- You are one auditor in a multi-agent tech-debt/critical-fix audit of the khelsutra-guru multi-repo workspace at C:/Users/avidu/Projects/khelsutra-guru. Today is 2026-07-02. All git repos were just pulled to current origin master/main.
- Repos: badminton-highlight-indexer (the ENGINE, Python, ~1500 tests, org Khelsutra), khelsutra (SPA web/ + marketing site/, TS/React, org avidullu), sports-data-collector, rally-annotator, sports-obsreport (small Python repos). Plus two NON-git dirs: "Annotation Setup", khelsutra-screencast.
- HARD RULES: READ-ONLY. Do NOT modify/create/delete any file, do NOT create branches/worktrees, do NOT pip install, do NOT run the full test suite (CI is trusted on master). You MAY run read-only git and gh commands (gh is authenticated), ruff/mypy IF already installed, and fast read-only scripts.
- EVERY claim must cite file:line (or a gh URL). A claim without file:line evidence is a hypothesis - mark it as such or drop it. The codebase drifts between sessions: re-verify anything you believe from docs against actual code.
- OWNER-GATED areas (flag as owner_gated=true, still report): alpha/serving go-live + Cloudflare dashboard work, Studio P10/P11 (corpus promotion + train/reeval), product/strategy/licensing decisions, real GPU/cloud spend, counsel ToS/DPA, repo rename off "badminton".
- Known-open items from project memory (VERIFY current state, don't trust blindly): engine PR #390 (D1 split main.py, open since 06-26) · audit doc docs/CODE_CLEANUP_AUDIT_2026-06-24.md remaining items D1, D13 (ByteTrack->trackers), Phase-3 (B3/B4/B6/B7/B8/D5/D6/D8-D10/03/04) · docs/DECISION_SNAPSHOT_REGRESSION.md next=P1 · docs/GOLDEN_REGRESSION_FIXTURES.md backlog M4/P0-rename/01/02 · engine issues #340 (Gemini re-bill), #332 (NSFW preflight), #349 (CPU-bound decode), #384 (CI single-source) · khelsutra issues #97-#114 UX/alpha backlog · DO droplet RETIRED 2026-06-25 (docs may still reference it; api.khelsutra.guru CNAME/CF-Access cleanup was pending).
- Categorize each finding: critical-fix (latent bug/correctness/data-loss/cost-leak), security, tech-debt (structure/duplication/dead code), test-gap, ci-infra, doc-rot, hygiene.
- Severity: P0 = latent bug or cost/data/security risk reachable in normal use; P1 = high-value debt blocking velocity or a stale-open thread; P2 = worthwhile cleanup; P3 = nice-to-have (report max 3 of these).
- Be selective: max ~12 findings, each genuinely worth acting on. Note tracked_in when an issue/PR/doc already tracks it (tracked is NOT done - still report). Your final message is data for an orchestrator, not prose for a human.

YOUR DIMENSION: cross-repo consistency + workspace hygiene across

C:/Users/avidu/Projects/khelsutra-guru.

Check: (1) engine->khelsutra API contract - the SPA (khelsutra/web/src) vs engine routes (backend/main.py, backend/api/) - any drift (endpoints, response fields like the new fps/resolution in engine #413 vs khelsutra #140 - did both sides land coherently?); (2) the vendored SPA bundle in the engine (backend/ui/) - how is it rebuilt, is there a documented + scripted sync process, is it currently stale vs khelsutra master (this is a recurring drift trap - assess); (3) duplicated conventions/CI boilerplate across the 5 repos (ruff/mypy configs, PR

templates, CLAUDE.md files - inconsistencies that cause friction); (4) workspace hygiene: for each repo run `git worktree list + git branch -a | wc -l + git status --porcelain - report stale worktrees/branches/dirty trees`; (5) each repo's CLAUDE.md accuracy spot-check. Cite file:line.

2. user (2026-07-02T02:47:22.592Z)

Annotation Setup
badminton-highlight-indexer
khelsutra
khelsutra-screencast
rally-annotator
sports-data-collector
sports-obsreport

3. user (2026-07-02T02:47:26.857Z)

```
=== badminton-highlight-indexer ===
C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer e410136 [master]
-- branches: 8
* master
  remotes/origin/HEAD -> origin/master
  remotes/origin/chore/teardown-accept-rally-vm-name
  remotes/origin/docs/nightly-person-cue-lane
  remotes/origin/docs/preserve-rally-detection-review
  remotes/origin/fix/rally-detection-edge-cases
  remotes/origin/master
  remotes/origin/refactor/p10-split-main
-- status:
-- head: e410136 api: /api/videos surfaces fps + resolution (from the stored ffprobe details)
(#413)
=== khelsutra ===
C:/Users/avidu/Projects/khelsutra-guru/khelsutra f9eea14 [master]
-- branches: 11
* master
  remotes/origin/chore/retire-droplet-parity
  remotes/origin/claude/badminton-academy-flyer-w3obpr
  remotes/origin/deploy/alpha-cloudflared-kit
  remotes/origin/docs/friendly-multilingual-project
  remotes/origin/docs/multilingual-expansion-proposal
  remotes/origin/feat/async-compile-poll
  remotes/origin/feat/flyer-nav-ownmodel-links
  remotes/origin/feat/link-academy-flyers
  remotes/origin/master
  remotes/origin/spa/api-base-env
-- status:
-- head: f9eea14 web: resilient video-summary readers + show fps in the library (#140)
=== sports-data-collector ===
C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector adcd206 [master]
-- branches: 9
  chore/reclassify-corpus-licenses
  docs/adr-002-annotation-proxy
  docs/compiler-padding-rationale
  feat/golden-corpus-n6
  fix/cvat-empty-corrupt-zip-gate
  fix/import-local-golden-label-guard
  fix/licensing-video-vs-annotation-gate
* master
  remotes/origin/master
-- status:
  M .gitignore
?? sports-data-collector/
-- head: adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9)
(#53)
=== rally-annotator ===
C:/Users/avidu/Projects/khelsutra-guru/rally-annotator 3d604bc [main]
```

```
-- branches: 6
* main
  remotes/origin/claude/intelligent-planck-fokdq2
  remotes/origin/feat/layout-snapshot-ci
  remotes/origin/feat/version-in-window-title
  remotes/origin/fix/gitattributes-eol-lf
  remotes/origin/main
-- status:
-- head: 3d604bc docs: sync READMEs with current state (multi-agent audit fixes) (#25)
=== sports-obsreport ===
C:/Users/avidu/Projects/khelsutra-guru/sports-obsreport 1f72c1c [main]
-- branches: 2
* main
  remotes/origin/main
-- status:
-- head: 1f72c1c feat: configurable highlights item_noun (v0.1.3)
```

4. user (2026-07-02T02:47:42.075Z)

warning: in the working copy of '.gitignore', LF will be replaced by CRLF the next time Git touches it

```
diff --git a/.gitignore b/.gitignore
index 52dac92..ac1420f 100644
```

```
--- a/.gitignore
+++ b/.gitignore
@@ -53,3 +53,7 @@ data/temp/
```

```
# Local DB backups (e.g. data/sports_metadata.backup-YYYY-MM-DD.db) ? never commit
data/*.backup-*.db
```

```
+
+# Model artifact bundles: weights are NEVER committed (private Releases); the
+# reviewable surface is each bundle's manifest.json (ADR-008).
```

```
+artifacts/
```

```
---nested dir---
```

```
README.md
```

```
cleanup.py
```

```
config
```

```
data
```

```
docs
```

```
pyproject.toml
```

```
src
```

```
tests
```

```
---local branches merged---
```

```
* master
```

```
---unmerged---
```

```
  chore/reclassify-corpus-licenses
```

```
  docs/adr-002-annotation-proxy
```

```
  docs/compiler-padding-rationale
```

```
  feat/golden-corpus-n6
```

```
  fix/cvat-empty-corrupt-zip-gate
```

```
  fix/import-local-golden-label-guard
```

```
  fix/licensing-video-vs-annotation-gate
```

```
---branch tips---
```

```
chore/reclassify-corpus-licenses: e5f799a chore(corpus): reclassify 23 ShuttleSet broadcast
videos as non-commercial (owner-approved) | ahead-of-master: 1
```

```
docs/adr-002-annotation-proxy: 265392c docs: ADR-002 - vendor annotation proxies generated in-
repo (curate prep-annotation-proxy) | ahead-of-master: 1
```

```
docs/compiler-padding-rationale: 4ce9e01 docs(cvat): update detect_issues padding rationale to
config-driven values | ahead-of-master: 1
```

```
feat/golden-corpus-n6: 2226e00 feat(corpus): ingest the n=6 owned golden corpus (262 rallies,
Proprietary/commercial) | ahead-of-master: 1
```

```
fix/cvat-empty-corrupt-zip-gate: 9a8937c fix(cvat): fail the delivery gate on
empty/corrupt/zipped CVAT instead of passing | ahead-of-master: 1
```

```
fix/import-local-golden-label-guard: b9a43a1 fix(import-local): protect golden labels + guard
the fractional-rally_number landmine | ahead-of-master: 1
```

fix/licensing-video-vs-annotation-gate: 9131b91 fix(licensing): gate commercial use on the VIDEO license, not the annotation license | ahead-of-master: 1

5. user (2026-07-02T02:47:54.786Z)

```
total 53
drwxr-xr-x 1 avidu 197609    0 Jun 19 21:38 .
drwxr-xr-x 1 avidu 197609    0 Jun 29 15:27 ..
drwxr-xr-x 1 avidu 197609    0 Jul  2 08:17 .git
-rw-r--r-- 1 avidu 197609  264 Jun  2 22:13 .gitattributes
drwxr-xr-x 1 avidu 197609    0 Jun 19 21:38 .github
-rw-r--r-- 1 avidu 197609 1139 Jun 19 21:38 .gitignore
-rw-r--r-- 1 avidu 197609 9583 Jun 19 21:38 README.md
-rwxr-xr-x 1 avidu 197609 6802 Jun 19 21:38 cleanup.py
drwxr-xr-x 1 avidu 197609    0 Jun  2 22:13 config
drwxr-xr-x 1 avidu 197609    0 Jun 19 21:38 data
drwxr-xr-x 1 avidu 197609    0 Jun 19 21:38 docs
-rw-r--r-- 1 avidu 197609 1211 Jun 19 21:38 pyproject.toml
drwxr-xr-x 1 avidu 197609    0 Jun 19 21:38 src
drwxr-xr-x 1 avidu 197609    0 Jun 19 21:38 tests
---is it a git repo?---
.
..
.git
.gitattributes
.github
.gitignore
README.md
cleanup.py
config
data
---diff vs parent src---
Only in /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/src/cli: __pycache__
Files /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/src/cli/curate.py and
/c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/sports-data-
collector/src/cli/curate.py differ
Files /c/Users/avidu/Projects/khelsutra-guru/sports-data-
collector/src/cli/normalize_annotations.py and /c/Users/avidu/Projects/khelsutra-guru/sports-
data-collector/sports-data-collector/src/cli/normalize_annotations.py differ
Only in /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/src/core: __pycache__
Files /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/src/core/db.py and
/c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/sports-data-
collector/src/core/db.py differ
Only in /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/src/core: manifests.py
Files /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/src/core/parsing.py and
/c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/sports-data-
collector/src/core/parsing.py differ
Only in /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/src/core/schemas:
__pycache__
Only in /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/src/core/schemas:
_mixins.py
Files /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/src/core/schemas/badminton.py
and /c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/sports-data-
collector/src/core/schemas/badminton.py differ
---cherry check (are branch commits content-equivalent in master?)---
chore/reclassify-corpus-licenses: - e5f799a019d9d57d7903a8cda3b1be259f06225d
docs/adr-002-annotation-proxy: - 265392c043ccb6717ada426422f8b8e5de36543c
docs/compiler-padding-rationale: - 4ce9e01ada9e74e0678232a5dd76e537370a0abd
feat/golden-corpus-n6: - 2226e0071fa1e9ab62a490927f0c83abcf6ed3c
fix/cvat-empty-corrupt-zip-gate: - 9a8937c767c1a2578f2266ff0c4ec821b0be55a4
fix/import-local-golden-label-guard: - b9a43a1e9e4f99d00bd762772ebf887bbc512d49
fix/licensing-video-vs-annotation-gate: - 9131b914bb8cbc3c0e4f0854a854c575f44d070c
```

6. user (2026-07-02T02:47:55.877Z)

53 MERGED feat/db-busy-timeout fix(db): SQLite busy_timeout so concurrent writers wait, not error (P9)
52 MERGED feat/import-annotations-fn feat(import): callable import_annotations + a maturity gate (P9)
51 MERGED feat/export-manifest-md5 feat(export): add source file MD5 to the eval/training manifest (P9)
50 MERGED fix/cvat-load-canonical-warn fix(cvat): warn on rows silently skipped by load_canonical_csv
49 MERGED docs/archive-code-health-plan-done docs: archive CODE_HEALTH_PLAN ? project DONE ?
48 MERGED refactor/a5-typed-manifest-models refactor(manifests): typed pydantic models for the export manifest (A5)
47 MERGED refactor/a2-curation-error refactor(curate): lift sys.exit out of _guard_import into CurationError (A2)
46 MERGED fix/a6-configurable-faq-path fix(reporting): make FAQ path configurable via env var (A6)
45 MERGED docs/next-steps-tracker-done docs: Tier D triage (D8) + NEXT_STEPS for parked A1/A3
44 MERGED refactor/t8b-curate-type-hints refactor(curate): concrete type hints for annotation handling (Tier C)
43 MERGED refactor/t4-config-driven-thresholds refactor(config): promote validation thresholds to config (Tier C)
42 MERGED fix/t7-defusedxml fix(rally_cvat): harden XML parsing with defusedxml (T7/OQ5)
41 MERGED refactor/t3-schema-versioning refactor(db): versioned schema migrations via PRAGMA user_version (T3)
40 MERGED refactor/t2-timestamp-mixin refactor(schemas): extract shared timestamp validator to mixin (T2)
39 MERGED fix/t9-t6-license-dedupe-strict-safeint refactor(validator,parsing): dedupe license whitelist + strict safe_int (T9/T6)
38 MERGED fix/b6-crawler-tempfile-cleanup fix(crawler): guarantee temp-file cleanup in finally (B6)
37 MERGED fix/b2-coachai-buffer-collision fix(coachai): log loudly on buffer-collision collapse (B2/OQ4)
36 MERGED fix/b3-boundary-grid-identical fix(rally_cvat): flag identical boundaries the grid check missed (B3)
35 MERGED fix/b1-is-commercial-validator fix(schemas): strict is_commercial ? license_type consistency-check (B1/D7)
34 MERGED fix/b5-b4-enum-crash-cricket-bool fix(db,curate): graceful enum rehydrate + cricket wicket truthiness (B5/B4)

7. user (2026-07-02T02:48:10.554Z)

031e410 feat(annotation): richer Roora schema - court + confidence + capture report + IAA (#29)
master
origin <https://github.com/avidullu/sports-data-collector.git> (fetch)
origin <https://github.com/avidullu/sports-data-collector.git> (push)

8. user (2026-07-02T02:48:11.638Z)

/c/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/backend:
__pycache__
annotations
api
assets
audio
billing.py
compute
config
detectors
eval
features
flywheel.py
infrastructure
main.py
pipeline
providers
reporting

results.py
sports
stitcher
storage
tools
ui
utils
worker

/c/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/backend/api:
__init__.py
__pycache__
auth.py
job_worker.py
static
uploads.py

/c/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/backend/ui:
THIRD_PARTY_NOTICES.md
assets
index.html
ui_version.json

9. user (2026-07-02T02:48:19.484Z)

```
1      {  
2        "spa_repo": "https://github.com/avidullu/khelsutra.git",  
3        "spa_commit": "ed33a515a4c68cb42134e8c7d4514df2addb9f54",  
4        "engine_api_version": 1,  
5        "content_hash": "b0801e7a231d4849cb33d374840f92d7ab2aca3f6ba9088124bbdd674ff15851",  
6        "built_at": "2026-06-22T17:59:59Z"  
7      }  
8
```

10. user (2026-07-02T02:48:21.016Z)

f9eea14 web: resilient video-summary readers + show fps in the library (#140)
9223ba9 docs: refresh khelsutra docs for today's storage-track work (P2-P5) (#139)
fb2f561 web: persistent "FFmpeg not detected" banner (hard-blocker warning) (#138)
8667dbf P5: SPA storage-medium picker (store a video into your medium) (#137)
cc06c32 feat(web): P8 annotate-in-UI ? AnnotatePanel over the vendored core (flagged) (#128)
9e9df72 feat(web): keep the video name visible after open (persistent current-video label) (#127)
926b53a feat(web): vendor the rally-annotator annotation core (P8 groundwork) (#126)
39c1349 feat(web): P6 ? "Your videos" library (reopen a stored video) (#124)
---current master sha---
f9eea14e05b5e3cc48e776757c38449d4cb5e211
2026-07-01 22:08:00 +0530

11. user (2026-07-02T02:48:30.074Z)

10
f9eea14 web: resilient video-summary readers + show fps in the library (#140)
9223ba9 docs: refresh khelsutra docs for today's storage-track work (P2-P5) (#139)
fb2f561 web: persistent "FFmpeg not detected" banner (hard-blocker warning) (#138)
8667dbf P5: SPA storage-medium picker (store a video into your medium) (#137)
cc06c32 feat(web): P8 annotate-in-UI ? AnnotatePanel over the vendored core (flagged) (#128)
9e9df72 feat(web): keep the video name visible after open (persistent current-video label) (#127)
926b53a feat(web): vendor the rally-annotator annotation core (P8 groundwork) (#126)
39c1349 feat(web): P6 ? "Your videos" library (reopen a stored video) (#124)
bd1cfa1 docs(e2e): make the screencast harness portable + turnkey across machines (#123)
d0285cd test(e2e): single-origin mock engine + Playwright screencast harness (#122)
---date of ed33a515---

ed33a51 2026-06-22 23:29:35 +0530 feat(build): poll the async /api/compile job for the reel (#329) (#108)

12. user (2026-07-02T02:48:31.169Z)

Found 22 files
docs\CODE_MAP.md
docs\NEXT_STEPS.md
backend\main.py
tests\test_ui_bundle.py
scripts\bundle_ui.sh
docs\data_pipeline\README.md
docs\data_pipeline\GOLDEN_SET_IMPLEMENTATION.md
docs\data_pipeline\GOLDEN_SET_PHASE1_VIDEOS.md
docs\archives\roora-vendor\ROORA_LABELING_GUIDELINES.md
docs\archives\roora-vendor\GOLDEN_SET_VENDORING.md
docs\archives\past_projects\DECOUPLED_COMPUTE\02-COMPUTE-MAP-AND-CONTRACT.md
docs\archives\decisions\DECISIONS.md
docs\RALLY_QUALITY_RESEARCH.md
docs\PACKAGING_PLAN.md
docs\MULTI_SIGNAL_FUSION_PLAN.md
docs\HOSTING_PLAN.md
docs\COMPUTE_DECOUPLED_SERVING\README.md
docs\BACKLOG.md
backend\ui\ui_version.json
backend\ui\assets\index-ukuS9wxa.js
backend\ui\THIRD_PARTY_NOTICES.md
.gitattributes

13. assistant (2026-07-02T02:48:40.711Z)

The vendored bundle is 10 web/ commits stale, including the #140 fps UI. Let me check the sync script and the bundle test.

14. user (2026-07-02T02:48:42.073Z)

```
1      #!/usr/bin/env bash
2      # Bundle the KhelSutra web SPA into the engine so a downloaded engine serves the real UI
3      # single-origin (PACKAGING_PLAN.md P1b).
4      #
5      # OPTION 1 (now, owner-chosen): build the SPA from a khelsutra checkout and vendor the
built `dist/`
6      # into `backend/ui/`, which is COMMITTED ? so the wheel + CI build fully
offline/reproducibly. Refresh
7      # the snapshot by re-running this on an SPA release; commit the diff.
8      #
9      # OPTION 3 (designed-for, future): point --from at a published dist tarball/URL (a
GitHub Release
10     # artifact) to vendor a pre-built SPA without a local khelsutra checkout or Node.
Already wired below.
11     #
12     # Usage:
13     #   scripts/bundle_ui.sh [SOURCE]
14     #   SOURCE = a khelsutra checkout dir (default: $KHELSUTRA_DIR or ../khelsutra)
[builds it]
15     #           | a path/URL ending .tgz/.tar.gz
[unpacks it]
16     #   BUILD=0 scripts/bundle_ui.sh <khelsutra-dir> # skip npm; reuse an existing
web/dist
17     set -euo pipefail
18
19     ENGINE_ROOT="$(cd "$(dirname "${BASH_SOURCE[0]}")/.." && pwd)"
20     UI_DEST="$ENGINE_ROOT/backend/ui"
21     SOURCE="${1:-${KHELSUTRA_DIR:-$ENGINE_ROOT/../../khelsutra}}"
```

```

22
23     stage_dist() { # populates $1 (a dist dir) into a clean $UI_DEST
24         local dist="$1"
25         [ -f "$dist/index.html" ] || { echo "ERROR: no index.html in $dist ? not a built SPA
dist" >&2; exit 1; }
26         rm -rf "$UI_DEST"
27         mkdir -p "$UI_DEST"
28         cp -r "$dist"/. "$UI_DEST"/
29     }
30
31     SPA_REPO="unknown"
32     SPA_COMMIT="unknown"
33
34     case "$SOURCE" in
35         *.tgz|*.tar.gz|http://*|https://*)
36             # OPTION 3: a published dist artifact.
37             tmp="$(mktemp -d)"; trap 'rm -rf "$tmp"' EXIT
38             if [[ "$SOURCE" == http*://* ]]; then
39                 echo "[bundle_ui] fetching dist artifact: $SOURCE"
40                 curl -fsSL "$SOURCE" -o "$tmp/dist.tgz"; SRC_TGZ="$tmp/dist.tgz";
SPA_REPO="$SOURCE"
41             else
42                 SRC_TGZ="$SOURCE"; SPA_REPO="$SOURCE"
43             fi
44             tar -xzf "$SRC_TGZ" -C "$tmp"
45             # accept either dist/* at the root or a top-level dir containing index.html
46             if [ -f "$tmp/index.html" ]; then DIST="$tmp"; else DIST="$(dirname "$(find "$tmp"
-name index.html | head -1)");" fi
47             stage_dist "$DIST"
48             ;;
49         *)
50             # OPTION 1: a khelsutra checkout ? build it.
51             WEB="$SOURCE/web"
52             [ -d "$WEB" ] || { echo "ERROR: $WEB not found (pass a khelsutra checkout dir or a
dist tarball)" >&2; exit 1; }
53             SPA_REPO="$(git -C "$SOURCE" config --get remote.origin.url 2>/dev/null || echo
"$SOURCE")"
54             SPA_COMMIT="$(git -C "$SOURCE" rev-parse HEAD 2>/dev/null || echo unknown)"
55             if [ "${BUILD:-1}" = "1" ]; then
56                 echo "[bundle_ui] building SPA in $WEB (npm ci && npm run build)"
57                 ( cd "$WEB" && npm ci && npm run build )
58             else
59                 echo "[bundle_ui] BUILD=0 ? reusing existing $WEB/dist"
60             fi
61             stage_dist "$WEB/dist"
62             ;;
63     esac
64
65     # Third-party notices for the vendored bundle (the python license-scan lane can't see
JS/font blobs).
66     cat > "$UI_DEST/THIRD_PARTY_NOTICES.md" <<NOTICE
67     # Third-party notices ? vendored KhelSutra SPA (`backend/ui/`)
68
69     This directory is the BUILT KhelSutra web SPA (MIT), vendored from $SPA_REPO @
$SPA_COMMIT by
70     `scripts/bundle_ui.sh` (PACKAGING_PLAN.md P1b). The minified `assets/*.js` bundle
includes
71     third-party runtime dependencies (React and the SPA's other libraries ? MIT/BSD), and
the
72     `assets/*.woff2` are Noto Sans Kannada/Devanagari under the SIL Open Font License 1.1
(decision D7).
73     Per the project's MIT/Apache/BSD-only guardrail, the SPA's dependencies are vetted in
the SPA repo;
74     see its package manifest/lockfile for the authoritative per-package licenses.
75     NOTICE

```

```

76
77     # Provenance + the SPA<->engine contract version (the engine's API_VERSION at bundle
time). The
78     # content hash is over the SPA build output only (relative paths ? reproducible across
checkouts),
79     # excluding the two generated metadata files.
80     ENGINE_API_VERSION="$(grep -E '^API_VERSION\b[^\=]*=' "$ENGINE_ROOT/backend/main.py" |
grep -oE '[0-9]+' | head -1)"
81     CONTENT_HASH="$( (cd "$UI_DEST" && find . -type f ! -name ui_version.json ! -name
THIRD_PARTY_NOTICES.md | LC_ALL=C sort | xargs sha256sum) | sha256sum | cut -d' ' -f1)"
82     BUILT_AT="$(date -u +%Y-%m-%dT%H:%M:%SZ 2>/dev/null || echo unknown)"
83
84     cat > "$UI_DEST/ui_version.json" <<JSON
85     {
86         "spa_repo": "$SPA_REPO",
87         "spa_commit": "$SPA_COMMIT",
88         "engine_api_version": ${ENGINE_API_VERSION:-null},
89         "content_hash": "$CONTENT_HASH",
90         "built_at": "$BUILT_AT"
91     }
92     JSON
93
94     echo "[bundle_ui] vendored SPA -> $UI_DEST (spa_commit=$SPA_COMMIT,
engine_api_version=$ENGINE_API_VERSION)"
95     echo "[bundle_ui] files:"; find "$UI_DEST" -type f | sed "s|$UI_DEST/||" | sort
96

```

15. user (2026-07-02T02:48:42.322Z)

```

1     """The bundled SPA is served single-origin + the SPA<->engine contract
(PACKAGING_PLAN.md P1b)."""
2
3     import json
4     import os
5
6     import pytest
7
8     pytest.importorskip("httpx")
9     from fastapi.testclient import TestClient
10
11     import backend.main as m
12
13     client = TestClient(m.app)
14     UI_DIR = os.path.join(os.path.dirname(os.path.abspath(m.__file__)), "ui")
15
16
17     def test_bundled_ui_snapshot_is_committed():
18         assert os.path.isfile(os.path.join(UI_DIR, "index.html")), (
19             "backend/ui/index.html missing ? the SPA snapshot must be committed (run
scripts/bundle_ui.sh)."
20         )
21         assert os.path.isfile(os.path.join(UI_DIR, "ui_version.json"))
22
23
24     def test_serves_spa_at_root():
25         r = client.get("/")
26         assert r.status_code == 200
27         # The bundled index.html references its single-origin /assets bundle (Vite
base='/').
28         assert "/assets/" in r.text
29
30
31     def test_serves_spa_assets_byte_for_byte():
32         assets = os.path.join(UI_DIR, "assets")
33         js = next(f for f in os.listdir(assets) if f.endswith(".js"))

```

```

34         r = client.get(f"/assets/{js}")
35         assert r.status_code == 200
36         with open(os.path.join(assets, js), "rb") as f:
37             assert (
38                 r.content == f.read()
39             ) # served verbatim (mimetype varies by OS, so compare bytes)
40
41
42     def test_api_version_reports_bundled_ui():
43         body = client.get("/api/version").json()
44         assert body["api_version"] == m.API_VERSION
45         assert body["ui"] is not None
46         assert body["ui"]["spa_commit"] # provenance recorded
47
48
49     def test_spa_engine_contract_version_matches():
50         """Skew guard: the committed snapshot records the engine API_VERSION it was built
51         against. If
52         someone bumps API_VERSION without re-running scripts/bundle_ui.sh, this FAILS ?
53         forcing a
54         re-bundle so the shipped SPA can't silently target a stale API contract."""
55         with open(os.path.join(UI_DIR, "ui_version.json"), encoding="utf-8") as f:
56             uiv = json.load(f)
57             assert uiv["engine_api_version"] == m.API_VERSION, (
58                 f"bundled SPA built against engine api_version {uiv['engine_api_version']} but
59                 engine is now "
60                 f"{m.API_VERSION} ? re-run scripts/bundle_ui.sh and commit backend/ui/."
61             )

```

16. user (2026-07-02T02:48:56.851Z)

```

269:     Includes `engine_api_version` ? the API_VERSION the snapshot was built against (skew
guard)."""
408:API_VERSION = 1
424:     "api_version": API_VERSION,
450:     "api_version": API_VERSION,
488:     """Pull `(fps, resolution)` from the persisted validation results, or `(None,
None)`. The
489:     `resolution_fps_check` validator already probed ffmpeg and stored
`details={width,height,fps}`
492:     if isinstance(r, dict) and r.get("validator_name") == "resolution_fps_check":
494:         w, h, fps = d.get("width"), d.get("height"), d.get("fps")
495:         resolution = f"{w}x{h}" if w and h else None
496:         return (fps if isinstance(fps, (int, float)) else None), resolution
505:     `created_at` (from `processed_at`), and `fps`/`resolution` (from the stored
ffmpeg
509:     fps, resolution = _video_dims(row.get("validation_results"))
516:     "fps": fps,
517:     "resolution": resolution,
1364:     # M9 (D9): edge-auth tenant resolution. DEFAULT-OFF ? owner_id None (the local single-
user,
1548:     """Shared compile resolution ? run at SUBMIT (fast-fail 4xx) AND in the WORKER (re-
resolve, since a
1739:     fps = cap.get(cv2.CAP_PROP_FPS)
1742:     start_frame = max(0, int((timestamp - 3.0) * fps))
1743:     end_frame = min(int(total_frames), int((timestamp + 3.0) * fps))
---routes---
backend/main.py:258:@app.get("/", response_class=HTMLResponse)
backend/main.py:379:@app.get("/api/config")
backend/main.py:391:@app.get("/api/health")
backend/main.py:419:@app.get("/api/version")
backend/main.py:429:@app.get("/api/capabilities")
backend/main.py:521:@app.get("/api/videos")
backend/main.py:537:@app.get("/api/reels")

```

```
backend/main.py:570:@app.get("/api/reels/{filename}")
backend/main.py:623:@app.get("/api/source/{video_id}")
backend/main.py:663:@app.post("/api/annotations")
backend/main.py:727:@app.post("/api/assets", status_code=201)
backend/main.py:1355:@app.post("/api/process", status_code=202)
backend/main.py:1431:@app.get("/api/jobs/{job_id}/cost")
backend/main.py:1446:@app.get("/api/videos/{video_id}/cost")
backend/main.py:1457:@app.get("/api/jobs/{job_id}")
backend/main.py:1492:@app.get("/api/highlights/{video_id}/{filename}")
backend/main.py:1523:@app.post("/api/query")
backend/main.py:1594:@app.post("/api/compile", status_code=202)
backend/api/uploads.py:70:@router.post("/api/upload/init")
backend/api/uploads.py:111:@router.put("/api/upload/{upload_id}/part/{index}")
backend/api/uploads.py:157:@router.post("/api/upload/{upload_id}/complete")
backend/api/uploads.py:188:@router.delete("/api/upload/{upload_id}")
```

17. user (2026-07-02T02:48:58.698Z)

183 /api/

18. user (2026-07-02T02:49:05.497Z)

```
20 /api/types
19 /api/process
19 /api/compile
11 /api/capabilities
10 /api/jobs/
10 /api/client
9 /api/videos
9 /api/query
7 /api/jobs/j
7 /api/assets
6 /api/jobs/{id}
6 /api/config
5 /api/annotations
4 /api/upload/up1/part/0
4 /api/upload/init
4 /api/upload/
3 /api/process.
3 /api/jobs/j1
2 /api/uploads.py
2 /api/jobs/{job_id}
2 /api/highlights/vid/reel.mp4
1 `/api/*`
1 /api/videos.
1 /api/upload/up1/part/
1 /api/upload/up1/complete
1 /api/upload/up1
1 /api/source/{video_id}
1 /api/source/abc
1 /api/source
1 /api/query.
1 /api/jobs/{job_id}.
1 /api/jobs/j4
1 /api/jobs/j3
1 /api/jobs/j2
1 /api/highlights/{video_id}/{filename}
1 /api/highlights/vidX/r.mp4
1 /api/highlights/vid/r.mp4
1 /api/highlights/abc123/my
1 /api/highlights/...
1 /api/compile.
1 /api/client.ts
```

19. user (2026-07-02T02:49:14.521Z)

```
web/src/api/types.ts
web/src/api/types.ts
web/src/components/VideoLibrary.test.tsx
web/src/lib/video.test.ts
web/src/lib/video.ts
---
web/src/api/types.ts:44:  fps?: number | null;
```

20. user (2026-07-02T02:49:20.768Z)

```
  play_segments: RallySegment[];
}

/** One row of the indexed-video library. GET /api/videos [backend/main.py:447] returns the most
 * recent first so the console survives a reload. Every field but `video_id` is
optional/defensive. */
export interface VideoSummary {
  /** The MD5 join key (the contract field). Optional defensively ? an older/regressed engine
may
  * send only the raw `id`; read it through `lib/video.ts::videoId`, never directly. */
  video_id?: string;
  /** Raw DB fields an engine may send instead of the enriched contract (resilience fallbacks).
*/
  id?: string;
  filepath?: string | null;
  match_name?: string | null;
  local_path?: string | null;
  fps?: number | null;
  resolution?: string | null;
  created_at?: string | null;
}

export interface VideosResponse {
  videos: VideoSummary[];
}

/** POST /api/annotations ? persist finished rally labels for a video as the annotator CSV (P8).
 * The engine writes the CSV verbatim (byte-compatible with the pipeline), keyed by the file
MD5. */
export interface AnnotationsResponse {
  status: string;
  num_rows?: number;
  csv_uri?: string;
  message?: string;
}

export interface CompileResponse {
  status: string;
  /** server-local path (NOT a URL) ? build the download URL with highlightUrl(). */
  filepath: string;
}

/** The terminal RESULT of an async reel compile (#329): the job's `result` payload, surfaced by
 * GET /api/jobs/{id} once the stitch finishes. `status` is the pipeline verdict ("success" |
 * "failed"); a "failed" result carries a `message`. The SPA builds the download URL with
=== lib/video.ts fps reader ===
2:// the engine SHOULD return the contract fields (`video_id`, a friendly `match_name`, `fps`,
3:// `resolution`), but if a build regresses to the raw DB shape (`id`, `filepath`, no name),
the library
32:/** A compact spec line, e.g. `1920x1080 · 30fps` (either part optional; "" when neither is
known). */
35: if (v.resolution) parts.push(v.resolution);
36: if (typeof v.fps === "number" && v.fps > 0) parts.push(`${Math.round(v.fps)}fps`);
```

21. user (2026-07-02T02:49:21.537Z)

```
    ``match_2026``). Works for a local path, a bucket key, or a ``gdrive://`` URI (rsplit on
    ``/``).
    Falls back to the whole ``filepath`` if there's no basename."""
    base = str(filepath or "").rstrip("/").rsplit("/", 1)[-1].rsplit("\\", 1)[-1]
    stem = os.path.splitext(base)[0]
    return stem or base or str(filepath or "")

def _video_dims(validation_results: Any) -> "tuple[Optional[float], Optional[str]]":
    """Pull ``(fps, resolution)`` from the persisted validation results, or ``(None, None)``.
    The
    ``resolution_fps_check`` validator already probed ffprobe and stored
    ``details={width,height,fps}``
    on its result (every branch) ? so the videos row carries it with no schema change / re-
    probe."""
    for r in validation_results or []:
        if isinstance(r, dict) and r.get("validator_name") == "resolution_fps_check":
            d = r.get("details") or {}
            w, h, fps = d.get("width"), d.get("height"), d.get("fps")
            resolution = f"{w}x{h}" if w and h else None
            return (fps if isinstance(fps, (int, float)) else None), resolution
    return None, None

def _video_summary(row: Dict[str, Any]) -> Dict[str, Any]:
    """Map a DB video row to the API/SPA contract. ADDITIVE ? keeps the raw ``id``/``filepath``
    (an
    existing test + any old client rely on them) and adds the fields the SPA library + the
    persistent
    "current video" label actually read: ``video_id`` (the MD5 join key = ``id``), a friendly
    ``match_name`` (so an uploaded video shows its filename, not a generic fallback),
    ``local_path``,
    ``created_at`` (from ``processed_at``), and ``fps``/``resolution`` (from the stored ffprobe
    validation details) so the library cards can show them. Mirrors the mock engine's shape so
    mock and
    real agree."""
    fp = row.get("filepath") or ""
    fps, resolution = _video_dims(row.get("validation_results"))
    return {
        **row,
        "video_id": row.get("id"),
        "match_name": _friendly_video_name(fp),
        "local_path": fp,
        "created_at": row.get("processed_at") or row.get("created_at"),
        "fps": fps,
        "resolution": resolution,
    }

@app.get("/api/videos")
def list_videos(request: Request, limit: Optional[int] = None):
    """List indexed videos (most recent first) so the console survives a reload ? previously
    only a
    single-row get existed (PACKAGING_PLAN.md P1). Rows are enriched to the SPA contract
    (``video_id`` + a friendly ``match_name`` derived from the filename) so the library shows
    real
    names and the persistent "current video" label survives an upload ? see
    ``_video_summary``."""
    return {"videos": [_video_summary(v) for v in
    request.app.state.db.list_videos(limit=limit)]}

def _reels_dir(cfg) -> str:
```

```
return getattr(getattr(cfg, "output", None), "output_dir", "output") or "output"
```

```
_REEL_EXTS = (".mp4", ".mov")
```

```
@app.get("/api/reels")
```

22. user (2026-07-02T02:49:37.978Z)

docs/PACKAGING_PLAN.md:17:3. ? **P2c first-run wizard ? DONE (2026-06-22).** 3-step wizard (`SetupWizard.tsx`) from `/api/capabilities`, 6 langs, shows once/browser ? khelsutra

[#88](https://github.com/avidullu/khelsutra/pull/88); **re-bundled into `backend/ui/`**

(`scripts/bundle\_ui.sh` @ khelsutra@`7ae9cee`) so `indexer --app` SERVES it ? engine

[#312](https://github.com/Khelsutra/badminton-highlight-indexer/pull/312) . Browser-verified e2e.

docs/PACKAGING_PLAN.md:67:Build `khelsutra/web` (`npm ci && npm run build`) **at package time** and vendor `web/dist/` into the wheel via a **hatchling force-include** ? a package dir (e.g. `backend/ui/`), served by FastAPI with an **SPA history-fallback** route. The SPA already calls **same-origin `BASE=/api`** `[verified khelsutra/web client.ts:7]` ? **zero SPA code change** for single-origin serving. Retire the legacy CWD-relative `backend/api/static` mount `[verified main.py:71-72]`; resolve paths via `importlib.resources`. `dist/` is gitignored ? the build must record a **bundle hash** + an `/api/version` handshake + a contract test pinning SPA?engine.

docs/PACKAGING_PLAN.md:108:| P1b | **Bundle the SPA single-origin**: source-agnostic `scripts/bundle\_ui.sh` (build `khelsutra/web`@pinned now; unpack a release tarball later ? option 3) ? committed `backend/ui/` snapshot (option 1, ~700KB) + `ui\_version.json` + `THIRD\_PARTY\_NOTICES.md`; hatchling auto-includes it (no force-include); FastAPI serves `/` + `/assets` (legacy fallback); `/api/version` reports the bundled UI; SPA?engine skew-guard contract test; `backend/ui/` -text`. | P1a | No (cross-repo) | ? |

[#302](https://github.com/Khelsutra/badminton-highlight-indexer/pull/302) |

docs/PACKAGING_PLAN.md:114:| **P2c** | **In-UI 3-step first-run wizard** (Gemini key / Compute / Videos) rendered from `/api/capabilities`; **6 langs** (en/kn/hi/es/da/id); guided key-setup (no key-write endpoint yet ? tells the user how to add `GEMINI\_API\_KEY`; offline `motion` needs none); shows once/browser. SPA `SetupWizard.tsx` = khelsutra

[#88](https://github.com/avidullu/khelsutra/pull/88) , then **re-bundled into `backend/ui/`** (P1b refresh) = engine

[#312](https://github.com/Khelsutra/badminton-highlight-indexer/pull/312) so `indexer --app` SERVES it. Browser-verified e2e. | P1b | No (cross-repo) | ? |

[#88](https://github.com/avidullu/khelsutra/pull/88) .

[#312](https://github.com/Khelsutra/badminton-highlight-indexer/pull/312) |

docs/PACKAGING_PLAN.md:171:- 2026-06-22 ? **P2c first-run wizard DONE + the loop CLOSED.** 3-step `SetupWizard.tsx` (from `/api/capabilities`, 6 langs, guided key-setup) merged on the SPA (khelsutra #88), then **re-bundled into `backend/ui/`** via

`scripts/bundle\_ui.sh`@khelsutra@`7ae9cee` (engine #312) so **`indexer --app` SERVES the wizard** (live-boot verified). \$7 P2c ?. **? v1 LIGHT is FEATURE-COMPLETE** ? only owner clean-box acceptance + ffmpeg-acquisition UX remain before the north-star (P3-cloud/P4).

23. user (2026-07-02T02:49:40.850Z)

```
=== badminton-highlight-indexer ===
ci.yml
nightly.yml
PULL_REQUEST_TEMPLATE.md
-- configs:
CLAUDE.md
mypy.ini
pyproject.toml
ruff.toml
=== khelsutra ===
test.yml
-- configs:
=== sports-data-collector ===
ci.yml
-- configs:
pyproject.toml
=== rally-annotator ===
```

```
ci.yml
web-ci.yml
pull_request_template.md
-- configs:
=== sports-obsreport ===
-- configs:
pyproject.toml
```

24. user (2026-07-02T02:49:53.232Z)

```
=== khelsutra CLAUDE.md search ===
=== sports-data-collector CLAUDE.md search ===
=== rally-annotator CLAUDE.md search ===
=== sports-obsreport CLAUDE.md search ===
=== sports-obsreport .github ===
ls: cannot access '/c/Users/avidu/Projects/khelsutra-guru/sports-obsreport/.github': No such
file or directory
LICENSE
README.md
pyproject.toml
schema
scripts
specs
src
tests
```

25. user (2026-07-02T02:49:55.594Z)

```
=== engine ruff.toml ===
```

Lint-only lane (no `ruff format` enforcement ? a whole-repo reformat would bury blame history for zero behavior value).

```
#
```

Selected: pyflakes + pycodestyle errors (ruff defaults) + pycodestyle warnings (W) + bugbear (B) + isort (I ? consistent import ordering).

```
target-version = "py38"
```

scratch/ = throwaway debugging scripts that may hit real APIs; it is also excluded from pytest collection (see pytest.ini).

```
extend-exclude = ["scratch"]
```

```
[lint]
```

```
extend-select = ["W", "B", "I"]
```

```
[lint.per-file-ignores]
```

Entrypoint scripts deliberately run setup (load_dotenv, logging config, sys.path) before importing backend modules ? import order is load-bearing.

```
"backend/main.py" = ["E402"]
```

```
"scripts/run_phase3_eval.py" = ["E402"]
```

```
"scripts/run_process_and_stitch.py" = ["E402"]
```

```
=== engine mypy.ini (head) ===
```

```
[mypy]
```

Lenient, GREEN baseline (quality sweep AI-4): the gate protects the clean modules and pins an explicit quarantine list below ? not a typing crusade.

Ratchet: when you touch a quarantined module, try removing its entry.

```
ignore_missing_imports = True
```

backend/ is a namespace package (no root __init__.py) ? required so mypy

maps files to backend.* module names instead of duplicating top-level names.

```
explicit_package_bases = True
```

--- Quarantine (known-loose modules, each with a reason) ---

Optional module globals (db_instance/global_config set by create_app() / main()).

Route handlers use request.app.state exclusively; helpers take explicit params.

backend.api.job_worker lifted from quarantine (now takes explicit config/db params).

```
[mypy-backend.main]
ignore_errors = True
```

Worker dispatcher: composes the SAME validator/segmenter/report building blocks a backend.main, which are typed `config: dict` but receive the AppConfig migration shim at runtime (the .get() compat layer). Shares backend.main's config-shim quarantine

```
=== sdc pyproject tool sections ===
=== obsreport pyproject tool sections ===
[tool.setuptools.packages.find]
where = ["src"]
```

```
[tool.pytest.ini_options]
pythonpath = ["src"]
testpaths = ["tests"]
```

26. user (2026-07-02T02:50:06.991Z)

```
=== sdc ci.yml ===
name: CI
```

```
on:
  push:
    branches: [master]
  pull_request:
```

```
jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-python@v5
        with:
          python-version: "3.11"
      - name: Install dependencies
        run: |
          python -m pip install --upgrade pip
          pip install -e ".[dev]"
      - name: Run tests
        run: python -m pytest -q
```

```
=== khelsutra test.yml ===
```

```
name: test
on:
  push:
    branches: [master]
  pull_request:
```

```
jobs:
  site:
    runs-on: ubuntu-latest
    defaults:
      run:
```

```

    working-directory: site
  steps:
    - uses: actions/checkout@v7
    - uses: actions/setup-node@v6
      with:
        node-version: "24" # node:sqlite is unflagged; server.test.js boots the VM host
        cache: "npm"
        cache-dependency-path: site/package-lock.json
    # jsdom is a TEST-ONLY devDep (the behavioral landing-page UX tests run the real inline
    # script in a DOM); the Cloudflare Worker + VM-host RUNTIMES stay dependency-free. `npm
test`
    # runs ALL paths: Worker+D1 logic (index.test.js), the booted VM host (server.test.js),
the
    # static content contract (assets.test.js), and the UX behaviour (ux.dom.test.js).
    - run: npm ci
    - run: npm test
    # Syntax-check the ops shell scripts (provision/teardown/smoke/parity) so a broken
    # script can't ship to the runbook. `bash -n` parses without executing.
    - name: shell scripts parse-check
      run: for f in scripts/*.sh; do echo "checking $f"; bash -n "$f"; done

web:
  # Operator-console SPA (web/): typecheck + unit tests (with a coverage floor) + production
build
  # gate the merge. `test:coverage` fails if coverage drops below the thresholds in
vite.config.ts,
  # so a merge can't silently regress test coverage.
runs-on: ubuntu-latest
defaults:
  run:
    working-directory: web
  steps:
    - uses: actions/checkout@v7
    - uses: actions/setup-node@v6
      with:
        node-version: "24"
        cache: "npm"
        cache-dependency-path: web/package-lock.json
    - run: npm ci
    - run: npm run typecheck
    - run: npm run test:coverage
    - run: npm run build

layout:
  # Real-browser layout-regression guard. The other suites run in jsdom, which has NO layout
engine
  # (getBoundingClientRect() is all zeros) and therefore cannot see header overflow / wrapping
/
  # clipping ? the responsive-nav bug that recurred. This drives a headless Chromium over the
real
  # pages at a spread of viewport widths x all 6 locales and asserts the header never
distorts.
runs-on: ubuntu-latest
defaults:
  run:
    working-directory: site
  steps:
    - uses: actions/checkout@v7
    - uses: actions/setup-node@v6
      with:
        node-version: "24"
        cache: "npm"
        cache-dependency-path: site/package-lock.json
    - run: npm ci
    - run: npx playwright install --with-deps chromium

```

```

- run: npm run test:layout
# Cross-locale render guard: drives the real pages in every locale and asserts the live
DOM
# matches the catalog (no raw keys / no shim regression / right lang / no console errors).
- run: npm run test:render
# Accessibility guard: axe-core WCAG 2.1 A/AA over every page × representative locales;
fails on
# any serious/critical violation (jsdom has no contrast/layout engine to catch these).
- run: npm run test:a11y

```

27. user (2026-07-02T02:50:17.119Z)

name: CI

on:

```

push:
  branches: [master]
pull_request:

```

permissions:

```

  contents: read

```

concurrency:

```

  group: ci-${{ github.ref }}
  cancel-in-progress: true

```

jobs:

```

# Lint gate (ruff.toml at repo root): pyflakes/pycodestyle + bugbear.
# Also the repo-wide syntax guard ? ruff reports invalid-syntax on every
# tracked .py, including scripts/ and training/, which the backend-only
# py_compile smoke sweep does not reach.

```

```

lint:
  name: Ruff lint
  runs-on: ubuntu-latest
  timeout-minutes: 10
  steps:
    - uses: actions/checkout@v6
    - uses: actions/setup-python@v6
      with:
        python-version: "3.13"
    - name: Install ruff
      run: python -m pip install ruff
    - name: Ruff check
      run: ruff check . --output-format github

```

```

# PII / attribution leak guard (tools/leak_scan.py): fails on a committed 32-hex MD5
# (the video_id shape) or a personal absolute user path in the shipped-code surface
# (backend/training/scripts/tools/eval_baselines; tests/ excluded ? legit fixtures). Real
# names are scanned only via a gitignored private list locally; CI runs the structural guard.

```

```

leak-scan:
  name: Leak scan (PII / attribution guard)
  runs-on: ubuntu-latest
  timeout-minutes: 10
  steps:
    - uses: actions/checkout@v6
    - uses: actions/setup-python@v6
      with:
        python-version: "3.13"
    - name: Leak scan
      run: python -m tools.leak_scan

```

```

# Internal-link guard (tools/check_md_links.py): fails on a broken internal markdown link
# (missing file or dead #anchor) in a LIVE doc. docs/archives/** is excluded ? frozen
# snapshots whose relative links are historical-by-design (see docs/DOC_STATUS.md §2).
link-check:

```

```

name: Markdown link check (live docs)
runs-on: ubuntu-latest
timeout-minutes: 10
steps:
  - uses: actions/checkout@v6
  - uses: actions/setup-python@v6
    with:
      python-version: "3.13"
  - name: Check internal markdown links
    run: python -m tools.check_md_links

# Type gate (mypy.ini at repo root): lenient green baseline with an explicit
# per-module quarantine list ? ratchet by removing entries as modules clean
# up. Typed core deps are installed so mypy sees their real signatures
# (without them ignore_missing_imports would Any-stub half the checks).
typecheck:
  name: Mypy (lenient baseline)
  runs-on: ubuntu-latest
  timeout-minutes: 10
  steps:
    - uses: actions/checkout@v6
    - uses: actions/setup-python@v6
      with:
        python-version: "3.13"
    - name: Install mypy + typed core deps
      run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-genai

numpy
  - name: Mypy

```

28. user (2026-07-02T02:50:20.508Z)

390 2026-06-26T14:28:06Z refactor/p10-split-main refactor(D1/P10): extract orchestration.py from main.py (god-module split, step 1)
 --- khelsutra open PRs ---

29. user (2026-07-02T02:50:31.352Z)

mergeable=CONFLICTING state=DIRTY updated=2026-06-26T14:37:18Z

30. user (2026-07-02T02:50:31.357Z)

```

1      # CLAUDE.md ? project memory / agent handoff
2
3      Local, AI-powered badminton (? racket-sports) highlight indexer + rally detector**:
4      FastAPI + SQLite + ffmpeg + GPU CV/AI, with paid Gemini as the cloud AI. Sibling
5      repos: `sports-data-collector`, a VLC-based `rally-annotator`.
6
7      ## Read these first
8      - docs/README.md ? doc index. docs/CODE_MAP.md ? code navigation + data
9        contracts (the cold-start map; read before touching code).
10     - Authoritative owned-model roadmap: docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md (+
the
11     live docs/NEXT_STEPS.md). These SUPERSEDE the 2026-06-07 strategy docs below where
they
12     disagree (licensing, Gemma-4 stance, base model, conversational-QA).
13     - Strategy (post-scaffolding): docs/COMMERCIALIZATION.md,
docs/PMF_MARKET_RESEARCH.md,
14     docs/PLATFORM_ARCHITECTURE.md, docs/DEFERRED_HARDENING_PLAN.md. Also
docs/ROADMAP.md
15     (offline-segmenter narrative) and docs/BACKLOG.md.
16     - Rally-detection quality track (design merged 2026-06-13, owner-gated, implementation
in progress):**
17     docs/HOW_RALLY_DETECTION_WORKS.md (2-layer mental model) ?
docs/MULTI_SIGNAL_FUSION_PLAN.md

```

```

18     (fuse shuttle + person + audio cues; the held-shuttle bug is the already-built
`rally_gate.py`
19     "Gate A" **not wired into production** ? cheapest win) ? `docs/EXPERIMENT_HARNESS.md`
(A/B + shadow
20     eval so cues land flag-gated/default-OFF with zero regression) ?
`docs/ROORA_LABELING_GUIDELINES.md`
21     (v8 vendor labeling spec). **`docs/NEXT_STEPS.md` has the prioritized pick-up for this
track.**
22     - **History:** `docs/archives/` ? ADRs (`decisions/DECISIONS.md`), research, and
23     **`checkpoints/`** (latest: `2026-06-strategy-foundation/`). **Archive policy:** only
move
24     **terminal-state** (DONE/REJECTED) work; live docs stay at `docs/` top level ? and
**before
25     archiving, HONESTLY update the doc first** (verify claims vs code ? flip status + a
completion
26     note ? update inbound refs ? keep the lineage ? then `git mv`), per the checklist +
living
27     lifecycle register in **`docs/DOC_STATUS.md`**.
28
29     ## Current state (June 2026)
30     - **Scaffolding complete** ; strategy foundation + owned-model roadmap merged. **No
31     platform/owned-model implementation has started.** A 2026-06-10 audit-driven hardening
32     sweep landed (licensing gate, detector keying, re-ingest safety, config, eval gate,
API
33     validation, reliability, CI + test coverage) ? see git history.
34     - **Two tracks, both owner-gated:** the **committed next PLATFORM slice = async job
model**
35     (`DEFERRED_HARDENING_PLAN.md` #16, single-tenant); the **owned-model track** starts at
a
36     greenlit Phase-0 milestone (`OWNED_MODEL_IMPLEMENTATION_PLAN.md`). **Do not start
37     implementation without explicit owner approval.**
38
39     ## Architecture in one breath
40     - **Pluggable registries** (ABC + `Registry` singleton): segmenters / providers /
validators
41     / sports / annotations ? see `backend/pipeline/segmenters/base.py`. *Future:*
42     `StorageBackend` + `ComputeBackend` registries (PLATFORM §3).
43     - `video_id` = MD5 of the file ? `backend/utils/hashing.py::compute_video_id` (one
helper).
44     - Typed config = `backend/config/` (Pydantic). DB schema evolves via the `PRAGMA
45     user_version` migration ladder in `backend/infrastructure/database.py`.
46
47     ## Committed guardrails (non-negotiable)
48     - **Paid LLMs (Gemini/OpenAI/Anthropic) = inference / serving / fallback ONLY ? never a
49     training data source** (terms/copyright). Commercial-clean licensing for EVERY
dependency ?
50     code, data, AND weights: **MIT / Apache-2.0 / BSD only** (ratified 2026-06-09, owner-
agreed).
51     - **Owned-model base:** **Qwen3-VL (Apache-2.0)** primary, InternVL3 (MIT) fallback;
**Gemma-4
52     is AMBER / bench-only** (Apache grant likely but Prohibited-Use posture unconfirmed ?
counsel
53     needed; do NOT build on it without sign-off). **Reject Gemma 3** (viral license).
**Exclude
54     minors** from the training pool; per-user training data needs a separate explicit opt-
in.
55     - **v1 = rally + highlight + history + correction loop** ; **gait/biometrics strictly
future**
56     (GDPR Art. 9).
57     - **Generic data schema** ; `coach ? athlete` is an **optional overlay** , not baked in.
58     - **Compute = GCP ONLY (ADR-013, 2026-06-20):** GPU training + serving run on **GCP**
(`deploy/gcp/` ;
59     Cloud Run+GPU for serving per ADR-012). **DigitalOcean is DROPPED for compute** (DO
GPU not enabled
60     on the account + not cleanly credit-funded; Paperspace is subscription-gated). Don't

```

```

re-evaluate
61     DO/Paperspace for compute without a new ADR; the `$200` DO credits go to **non-compute
only**. The
62     `deploy/digitalocean/*` scripts are kept solely as provider-agnostic Linux-GPU setup
reused by GCP.
63
64     ## Workflow conventions
65     - **Branch off `master`; ONE PR per work item; NO PR stacking.** Don't open a PR unless
asked.
66     - Tests: `python run_all_tests.py` (offline, fully mocked). Python 3.10+ (the `trackers`
package requires ?3.10); `ffmpeg` required.
67     - Never push to the default branch directly; descriptive commit messages.
68     - **Substantial design/project work gets a tracked project doc** following
`docs/PROJECT_DOC_TEMPLATE.md` (status header · decisions-locked · progress tracker ·
definition-of-done; archived on completion). Precedents: `docs/RALLY_QUALITY_RESEARCH.md` §7,
`docs/GOLDEN_REGRESSION_FIXTURES.md`.
69     - **Launch / "bit-flip" PRs** (a default-flip that drives a tracked project to terminal
state) MUST, in the SAME PR, run the `docs/DOC_STATUS.md` §2 archival due-diligence on that
project's docs (verify-vs-code ? flip status + completion note ? update inbound refs ? keep ADR
lineage ? `git mv` to `docs/archives/` ? record in DOC_STATUS §3), and fix any internal-link
breakage the move causes. Enforced by `.github/PULL_REQUEST_TEMPLATE.md` + the `link-check` CI
job (`python -m tools.check_md_links`).
70     - ?? Hand the heaviest machine-side checks (real video, the full GPU/WSL/WASB suite) to
the owner,
71     who has the golden corpus + labelled videos; verify pure-logic paths locally.
72

```

31. user (2026-07-02T02:50:55.544Z)

```

docs/HOSTING_PLAN.md
docs/NEXT_STEPS.md
docs/SETUP_NEW_MACHINE.md
docs/UI_ALPHA_EXECUTION_PLAN.md
docs/UI_PROPOSAL.md

```

32. user (2026-07-02T02:52:16.916Z)

```

deploy/DEPLOY_ALPHA.md
deploy/droplet/README.md
deploy/GPU_VM_SERVING.md
deploy/stage_from_bucket.sh
docs/DRY_RUN_CUJ.md
docs/LOCAL_APPLIANCE_ARCHITECTURE.md
docs/PORTABILITY.md
docs/README.md
NEXT_STEPS.md
site/README.md
site/scripts/droplet-smoke.sh
site/scripts/provision-vm.sh
site/scripts/teardown-vm.sh
site/src/server.test.js
web/src/api/client.test.ts
web/src/api/client.ts
web/src/vite-env.d.ts

```

33. user (2026-07-02T02:52:25.904Z)

```

12:dependency-free **Node + SQLite** host (`site/server.js`), **live + browsable on the DO
droplet ?
14:6-language localization, the per-rally MVP, droplet portability). Docs:
61:- **Droplet ? production** ? TLS + a subdomain (`vm.khelsutra.guru` via Caddy) if it should
be a real
64: (`deploy-test@`, `email-test@`). Revoke the droplet SSH key (`~/ssh/khelsutra_droplet`)
when done.

```

77: profile of the engine's decoupled `StorageBackend` + `backend.worker` architecture. MVP runs on the CPU droplet via
78: ****Gemini/stub****; ****real local-GPU detection needs a GPU host**** (DO droplet / GCE VM / Cloud Run+GPU ? owner priority
=== DEPLOY_ALPHA head ===

Deploy KhelSutra Studio for alpha ? a shareable page, no SSH

> ****Status:**** `LIVE` (deploy live-verified on the droplet 2026-06-22 ?
`deploy/droplet/README.md`; only a rotated Gemini key + a browser OTP login remain to exercise the upload?reel CUJ) . ****Topology:**** split ? SPA on Cloudflare Pages + engine on an `api.` cloudflared tunnel, Cloudflare Access over both (the engine's `HOSTING\_PLAN.md` \$2 alpha plan).
 . ****Cost:**** ~\$0/mo standing (Pages + Tunnel + Access are free ?50 users) + Gemini per-run (~\$0.05/match).

This turns the local-only flow (run engine + SPA on your box, drive it yourself) into a ****page** you can share with alpha testers** ? they open `app.khelsutra.guru`, sign in (email allowlist), point at a video, and get a reel. No SSH, no home IP exposed, no inbound ports.

> ****Outgrown the droplet?**** To move the live engine onto a ****GCP L4 GPU VM**** (more CPU/RAM for ffmpeg
> + optional WASB GPU detection) and cut the tunnel over to it, see
****GPU_VM_SERVING.md****
> ? same tunnel + Access, no DNS change, droplet kept as fallback (~\$500?600/mo while the VM runs).

0. The shape

```
...
tester browser
  ? https
  ?
Cloudflare edge (free): DNS . TLS . Access (email-OTP allowlist) in front of BOTH hostnames
  ??? app.khelsutra.guru   ? Cloudflare Pages           (the web/ SPA; built with VITE_API_BASE)
  ??? api.khelsutra.guru   ? cloudflared Tunnel (outbound-only) ?? 127.0.0.1:8000 on YOUR box
                                                                    ? FastAPI engine
                                                                    ? (Gemini segmenter; GPU
optional)
                                                                    ? async jobs, SQLite,
output/ on disk
...
```

- ****No inbound ports.**** `cloudflared` dials out; the engine stays bound to `127.0.0.1:8000`. The box's home/public IP is never exposed.
- ****Auth at the edge.**** Cloudflare Access gates both hostnames; the engine *also* verifies the `Cf-Access` JWT in-app (defence in depth ? `security.edge\_auth\_enabled`), so a direct hit to the tunnel can't spoof identity.
- ****Cross-origin, on purpose.**** `app.` (SPA) and `api.` (engine) are different origins ? the SPA is built with `VITE\_API\_BASE=https://api.khelsutra.guru/api` (PR #64) and the engine locks CORS to the SPA origin via `RALLY\_CORS\_ALLOW\_ORIGINS` (no engine code change ? it's already env-driven).

> ****Alternative (single-origin):**** if you'd rather run ****one box**** that serves the SPA *and* the API behind one origin (Caddy reverse-proxy, no Pages, no cross-origin), see
`LOCAL\_APPLIANCE\_ARCHITECTURE.md` D1. This runbook ships the ****split**** topology you chose (stable Pages-hosted page + a thin API tunnel).

=== client.ts droplet ref ===

25: * (e.g. `https://api.khelsutra.guru/api`), and lock the engine's `RALLY\_CORS\_ALLOW\_ORIGINS` to the

34. user (2026-07-02T02:52:41.879Z)

---retire-droplet-parity branch---

e9ae3c2 chore(ci): retire the live-parity monitor ? droplet mirror decommissioned

```
70083ad docs(deploy): fix $9 teardown command ? it no-ops (keeps billing) without RALLY_GPU_VM + zone (#116)
8316798 feat(deploy): stage_from_bucket.sh ? process bucket footage on the L4 VM without a browser upload (#115)
117 MERGED chore(ci): retire the live-parity monitor ? droplet mirror decommissioned
```

35. user (2026-07-02T02:52:55.866Z)

```
NOT in master (squash?)
merged=2026-06-25T11:20:58Z files=.github/workflows/parity.yml, site/scripts/parity-check.sh
---does master still have the parity workflow?---
test.yml
```

36. user (2026-07-02T02:53:10.087Z)

```
docs/NEXT_STEPS.md:33:15. Owner housekeeping: retire the shared **DO droplet** + **DO token**; collector **md5** export (collector repo).
docs/HOSTING_PLAN.md:78:| **All-DigitalOcean** | ~$11 ? ~$17 ($6 droplet + Spaces $5: 250 GiB + **1 TiB egress included**) | One simple vendor; Spaces' included egress is the best video-download deal; per-second billing | No credits; no scale-to-zero; still needs a tunnel-equivalent to reach the box (run `frp`/WireGuard on the droplet) |
```

37. user (2026-07-02T02:53:37.397Z)

```
/** motion-only fabricated sub-events ? DO NOT surface. */
events?: unknown[];
}
web/src/api/client.ts:87: * processing. POST /api/process {video_path} ? 202 {job_id, ...} [main.py:260]. Fast-fails 400
web/src/api/client.ts:105:/** Poll a submitted job. GET /api/jobs/{job_id} [main.py:303]. 404 ? unknown job_id. */
web/src/api/client.ts:110:/** List indexed videos, newest first ? the "my videos" library. GET /api/videos [main.py:447].
web/src/api/client.ts:156:/** List a video's detected rallies. POST /api/query {video_id} ? {play_segments} [main.py:331]. */
web/src/api/client.ts:223:/** The stream/download URL for a compiled reel. GET /api/highlights/{video_id}/{filename} [main.py:307]. */
web/src/api/client.ts:228:/** Read engine config (for the stitching floor). GET /api/config [main.py:121]. */
web/src/api/client.ts:238:/** The shot-count floor /api/compile silently enforces ? the UI must warn, not hide [main.py:362]. */
web/src/api/types.ts:3:// backend/main.py:331-396 (query/compile), :307-328 (highlights), :121 (config).
web/src/api/types.ts:33:/** One row of the indexed-video library. GET /api/videos [backend/main.py:447] returns the most
web/src/api/types.ts:143:/** POST /api/process ? 202 [backend/main.py:260]. The job is queued; poll GET /api/jobs/{job_id}. */
```

38. user (2026-07-02T02:53:39.961Z)

```
Archive
Golden Labelled
Trajectories
golden_labels_badminton.csv
rally_review.csv
---screencast---
b-205.png
b-218.png
b-228.png
cuj-1-ingest.png
cuj-2-uri.png
cuj-3-processing.png
cuj-4-rallies.png
cuj-5-reel-ready.png
```

cuji-downloaded-reel.mp4
cuji-screencast.mjs
deep-test.mjs
esdaid-screencast.mjs
eval-after
eval-before
eval-shots.mjs
4.9G /c/Users/avidu/Projects/khelsutra-guru/Annotation Setup
107M /c/Users/avidu/Projects/khelsutra-guru/khelsutra-screencast

39. user (2026-07-02T02:53:54.863Z)

---annotation setup csv sizes---
-rw-r--r-- 1 avidu 197609 2508 Jun 6 14:16 /c/Users/avidu/Projects/khelsutra-guru/Annotation Setup/golden_labels_badminton.csv
-rw-r--r-- 1 avidu 197609 1377 Jun 6 13:14 /c/Users/avidu/Projects/khelsutra-guru/Annotation Setup/rally_review.csv
Adarsh and Avi.rallies.csv
KushagraYashAviFirstVideo.rallies.csv
LargeTestVideo.rallies.csv
MahadevpuraSingles.rallies.csv
TestLargeVideo.rallies.csv
gBaaddy20May2026.rallies.csv

40. user (2026-07-02T02:54:05.959Z)

9e6f204 feat(corpus): ingest the n=6 owned golden corpus (262 rallies, Proprietary/commercial) (#19)
9e6f204 feat(corpus): ingest the n=6 owned golden corpus (262 rallies, Proprietary/commercial) (#19)
data/sports_metadata.db | Bin 200704 -> 233472 bytes
1 file changed, 0 insertions(+), 0 deletions(-)
---tracked data files---
data/mock_local_badminton_anno.json
data/roora_badminton_canonical.csv
data/roora_badminton_issues.txt
data/sports_metadata.db
data/staging_coachai_match.csv
data/validation/pilot_badminton_01_report.json
data/validation/pilot_badminton_01_report.md
data/validation/pilot_badminton_01_review.csv

41. user (2026-07-02T02:54:31.021Z)

engine chore/teardown-accept-rally-vm-name: 392 MERGED 351 MERGED
engine docs/nightly-person-cue-lane: 389 MERGED
engine docs/preserve-rally-detection-review: 393 CLOSED
engine fix/rally-detection-edge-cases: 394 MERGED
khelsutra chore/retire-droplet-parity: 117 MERGED
khelsutra claude/badminton-academy-flyer-w3obpr: 40 MERGED
khelsutra deploy/alpha-cloudflared-kit: 65 MERGED
khelsutra docs/friendly-multilingual-project: 41 MERGED
khelsutra docs/multilingual-expansion-proposal: 58 MERGED
khelsutra feat/async-compile-poll: 108 MERGED
khelsutra feat/flyer-nav-ownmodel-links: 44 MERGED
khelsutra feat/link-academy-flyers: 43 MERGED
khelsutra spa/api-base-env: 64 MERGED

42. user (2026-07-02T02:54:44.383Z)

OPEN ci: single-source the shared offline guard lanes + coverage floor via a workflow_call reusable
388 refactor: migrate async-job plumbing (_get_executor/_drain_due_jobs/_arm_wake_timer) out of main.py

384 ci: single-source the shared offline guard lanes + coverage floor via a workflow_call reusable
349 WASB native: decode is CPU-bound & single-threaded in streaming mode ? L4 GPU ~80% idle (no NVDEC/parallel decode)
340 [cost] /api/process re-runs the paid segmenter for an already-processed video ? refresh/"Try again" re-bills Gemini
332 [safety] Content-safety pre-flight litmus: flag non-sports / NSFW video BEFORE sending it to Gemini
331 [perf] Hosted processing is CPU-bound (ffmpeg/cv2) ? measured per-phase breakdown + levers (GPU/NVENC, parallel extraction, async compile)
330 [optimization] MD5 segmentation reuse ? skip Gemini re-segmentation when the same video (video_id) is re-uploaded
328 Surface the per-video runlog / execution trace to the user (endpoint + UI), not just server logs
327 Allow gs:// / s3:// / gdrive-api:// ingest on an exposed (jailed) box ? scheme-exempt recognized storage backends from the path-jail
326 Hosted-app observability: request tracing + replay for debugging production CUJ failures
301 T11 follow-up: run the upload retention sweep as an automated scheduled job (hosted alpha)
175 Shared cloud container (OneDrive) for GPU-extractor golden data ? sync GPU box ? dev sessions
172 Calibrate the ship-gate F1@tIoU target (currently provisional ? NOT owner-ratified)

43. user (2026-07-02T02:55:14.688Z)

DECISIONS.md
INGESTION_PIPELINE.md
MANUAL_INGESTION_POSITIONING.md
NEXT_STEPS.md
PIPELINE_ROLE.md
annotation
archives

44. user (2026-07-02T02:55:54.059Z)

/c/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/docs/archives/decisions/DECISIONS.md:388:## ADR-008: Owned-model topology ? artifact boundary now, repo split at Gen-2
/c/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/docs/archives/decisions/DECISIONS.md:498: failure mode is documented in ADR-008's rationale) with pointer READMEs in `khelsutra`.
/c/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/docs/archives/decisions/DECISIONS.md:691:`ADR-005` (permissive licensing ? engine + image stay MIT/Apache/BSD) · `ADR-008` (train==serve featurizer single-sourcing ? the parity discipline this extends to local?cloud) · `ADR-009` (**local-first delivery** ? origin of the truly-local profile) · `ADR-010` (DECOUPLED D0?GCP) · `ADR-011` (**deferred M3 serving + M4 billing** ; overrode "rent nothing / not a service"; \$06 gates = DPA/consent, collector `md5`, WASB-C3) · `ADR-012` (**Cloud Run+GPU = serving model** ; NVDEC lever) · `ADR-013` (**GCP sole compute stack**) · **? ADR-014** (this). Non-ADR: realizes `PLATFORM\_ARCHITECTURE.md`'s `ComputeBackend`/`StorageBackend` registries; the truly-local profile = `VIDEO\_LOCALITY\_MODEL.md` L0 + `DESKTOP\_APP\_PLAN.md` `LocalDesktop`; informed by the archived `DEPLOY\_REPORT\_GCP\_2026-06-20` (cross-GPU sub-pixel parity finding).
/c/Users/avidu/Projects/khelsutra-guru/sports-data-collector/.gitignore:58:# reviewable surface is each bundle's manifest.json (ADR-008).

45. user (2026-07-02T02:56:03.531Z)

CLAUDE.md
README.md
__pycache__
artifacts
backend
backend.db
config.json
deploy

docs
eval_baselines
indexer.db
mypy.ini
output
pyproject.toml
pytest.ini
rally_index.db
requirements.txt
ruff.toml
run_all_tests.py
run_eval.py
run_regression.py
scratch
scripts
tests
tools
training
README.md
__init__.py
__pycache__
gen0
requirements-train.txt

46. assistant (2026-07-02T02:56:56.787Z)

All checks complete. Compiling findings.

47. user (2026-07-02T02:58:05.059Z)

Structured output provided successfully