

claude-windows / c1fcb4aa-537d-41fd-858e-53f93349bf5a

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\subagents\workflows\wf_7049e1ed-b93\agent-a8227855a44240c6c.jsonl`  
- SHA-256: `3a918bcd1fd060db79c2935b4b458895d63c61269b1c13970f40896fafbaa7ae`  
- Source modified: `2026-07-02T04:58:28+00:00`  
- Imported at: `2026-07-05T16:48:24+00:00`  
- project: `wf_7049e1ed-b93`  
- session_id: `c1fcb4aa-537d-41fd-858e-53f93349bf5a`
```

Transcript

1. user (2026-07-02T04:45:39.418Z)

You are a senior engineer producing ONE isolated PR. Today is 2026-07-02.

REPO: C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer (a shared checkout ? follow the safety rules exactly).

AUDIT FINDINGS YOU ARE FIXING: R1-17 (scratch never reclaimed) + R1-03 overlap ? Grep C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/AUDIT_REPORT_khelsutra-guru_2026-07-02.md for these IDs and read the full evidence + verifier notes before coding.
TASK BRIEF:

In C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer: reclaim remote-input scratch. (1) backend/storage/__init__.py:89 acquire_local_input mkdtemp-per-call (rally-input-*): give worker job paths a per-job scope that deletes the scratch dir in a finally (the worker owns the job lifecycle ? check backend/worker/__main__.py + backend/api/job_worker.py call sites); (2) backend/compute/cloud_run.py _read_json marker reads: use tempfile.TemporaryDirectory context; (3) add an age-based startup sweep for orphaned rally-input-*/rally-worker-* dirs older than N days (default 3) in the server boot path, default-ON but logging what it removes, with a config kill-switch. DO NOT touch the /api/source path (a sibling PR adds a keyed cache there ? avoid backend/main.py get_source entirely to prevent conflicts). Tests: finally-cleanup on success+failure, sweep age-threshold respects the kill-switch.

BRANCHING SAFETY (mandatory ? shared checkout, sibling sessions exist):

1. In the MAIN repo checkout run: git fetch origin
2. Create an ISOLATED WORKTREE: git worktree add "C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-<shortname>" -b <branch-name> origin/master (branch explicitly FROM origin/master ? never from the current local branch; use origin/main for repos whose default is main)
3. Do ALL work inside the worktree. Before committing: git branch --show-current + git log --oneline -1 must show YOUR branch on the origin-master base.
4. Stage EXPLICIT paths only (git add <file> <file>) ? NEVER git add -A / . / -u
5. Before opening the PR: git log --oneline origin/master..HEAD must list ONLY your commits.

COMMIT/PR CONVENTIONS:

- Commit message: conventional style matching the repo's history; end the message with a blank line then: Co-Authored-By: Claude Fable 5 <noreply@anthropic.com>
- PR body: what/why, a "## Verification" section stating exactly what you ran locally and what CI covers, and end with: ? Generated with [Claude Code](https://claude.com/claude-code)
- Do NOT merge the PR. Do NOT remove the worktree (a fixer may need it). Return the PR number, branch, worktree path.

QUALITY BAR:

- Read the repo's CLAUDE.md first; for the engine also docs/CODE_MAP.md \$0 (file placement) before adding any new file.
- Add/extend tests for every behavior change. Run the targeted tests you touched (python -m pytest tests/<relevant> -q) + ruff check on changed files; if mypy config exists, run mypy on touched modules. Full suite is CI's job ? do not run it.

- Default-safe: new knobs default OFF / behavior-preserving unless the task explicitly says otherwise.
- The full audit evidence is in C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/AUDIT_REPORT_khelsutra-guru_2026-07-02.md ? Grep it for your finding IDs and READ the verifier notes; they contain corrections (exact file:line, scoping traps) you MUST honor.

If the brief's premise turns out false against current code (something already fixed, file moved), adapt honestly and say so in notes ? never force a pointless diff. If you genuinely cannot complete (env broken, auth missing), open NO PR and explain in notes with pr_number=0.

2. user (2026-07-02T04:45:47.186Z)

70:### R1-03 [P1/M] critical-fix ? badminton-highlight-indexer
196:### R1-17 [P2/M] tech-debt ? badminton-highlight-indexer

3. user (2026-07-02T04:45:48.260Z)

CLAUDE.md
README.md
__pycache__
artifacts
backend
backend.db
config.json
deploy
docs
eval_baselines
indexer.db
mypy.ini
output
pyproject.toml
pytest.ini
rally_index.db
requirements.txt
ruff.toml
run_all_tests.py
run_eval.py
run_regression.py
scratch
scripts
tests
tools
training

4. user (2026-07-02T04:45:56.298Z)

55 - evidence: gh pr view 390: state=OPEN, mergeable=CONFLICTING, mergeStateStatus=DIRTY, updatedAt=2026-06-26T14:37Z, 3 files (backend/main.py, backend/orchestration.py, mypy.ini). Since then >=8 PRs modified backend/main.py (#400, #405-#409, #412, #413 per git log --since=2026-06-25 -- backend/main.py). wc -l backend/main.py = 2198 vs the audit doc's '1448 lines' claim (docs/CODE_CLEANUP_AUDIT_2026-06-24.md:279); 18 @app. routes and 45 print() calls now live in it; backend.main still mypy-quarantined (mypy.ini:15). No backend/orchestration.py exists on master.

56 - why it matters: This is the last HIGH-severity item of the cleanup audit (P10, the only non-Phase-3 ? row in its \$6 tracker) and it is compounding: every new endpoint PR lands in the god-module, widening the conflict surface and making the split more expensive each week. The 3-file extraction in #390 is now unmergeable.

57 - proposed fix: Close #390 and re-cut the orchestration.py extraction fresh from origin/master (small, ~1 session), or rebase it now. Then continue the split (CLI handlers -> backend/cli/, new /api/source|api/assets|api/capabilities routes -> backend/api/) before more endpoints accrete. Update the audit doc's stale '~1450 lines' sizing while at it.

58 - verifier notes: Minor sharpenings only: (a) git log --since=2026-06-25 --

backend/main.py actually shows 10 commits (adds #387 repo-wide ruff format ? itself a likely source of #390's conflicts ? and #380), so ">=8 PRs" is understated; (b) the mypy quarantine block is mypy.ini:15-16 ([mypy-backend.main] header at 15, ignore_errors at 16); (c) the audit doc's "1448 lines" statement is in the SChangelog 2026-06-25 "Accuracy fixes" bullet (~line 278), and the stale "~1450" sizing also appears at lines 69, 198, and 247 ? all four spots need the size refresh the fix proposes.

59 - verifier reasoning: Every load-bearing claim reproduced. (1) `gh pr view 390` returns state=OPEN, mergeable=CONFLICTING, mergeStateStatus=DIRTY, updatedAt=2026-06-26T14:37:18Z, exactly 3 files: backend/main.py (+19/-595), backend/orchestration.py (ADDED, +635), mypy.ini (+6). (2) Current working tree (branch master, HEAD e410136 = PR #413): `wc -l backend/main.py` = 2198; 1448?2198 is +51.8% ? the claimed 52%; the 1448 baseline is documented in docs/CODE_CLEANUP_AUDIT_2026-06-24.md changelog ("main.py size '~1300' ? '~1450' (now 1448 lines)", 2026-06-25 entry, ~line 278-279). (3) Not fixed: backend/orchestration.py does not exist in the tree, and git log --since=2026-06-25 -- backend/main.py shows the file only grew ? commits for PRs #413, #412, #409, #408, #407, #406, #405, #400 (the claimed 8; the log actually lists 10 incl. #387 ruff-format and #380). (4) grep counts on backend/main.py: 18 lines matching `^@app\.` and 45 `print(` occurrences ? exactly as claimed. (5) mypy.ini:15-16 = `[mypy-backend.main]` / `ignore\_errors = True` (quarantine intact). (6) Audit doc \$6 tracker line 198: P10 row "D1: Split main.py god-module (~1450 lines ? smaller modules)" is ?; the only other ? row is line 205 (P16+ Phase 3 bucket), so P10 is indeed the sole non-Phase-3 open item; D1 is rated HIGH at doc line 69. Severity P1 is honest for a tech-debt audit: the concrete failure is already materialized (the 635-line extraction in #390 is unmergeable today), and it compounds ? 8 endpoint-adding PRs landed in the god-module in the week since the PR was cut, all inside a mypy ignore_errors module, so each week of delay increases rework cost and unchecked-type surface. Not owner-gated: internal refactor, no serving go-live / corpus-promotion / spend / licensing implications ("P10" here is the audit tracker row, not Studio P10). Proposed fix (close-and-recut or rebase from origin/master, then continue the split per the doc's own plan at line 247, and fix the stale ~1450 sizing) is sane and correctly scoped; effort M is right given the fix includes continuing the split beyond the re-cut.

60
61 ### R1-02 [P1/M] critical-fix ? badminton-highlight-indexer
62 **/api/process re-runs the paid Gemini segmenter for an already-processed video (re-bill on refresh/'Try again') ? still unfixed at HEAD**
63 - anchor: `backend/main.py:1172` · tracked in: Khelsutra/badminton-highlight-indexer#340 (also #330) · finder: engine:correctness
64 - evidence: backend/main.py:1171-1172 computes video_id and `was\_reingest` = db.get_video(video_id) is not None`, but was_reingest is only used for the report/destroy-after-confirm; there is NO reuse/skip branch ? the segmenter always runs (main.py:1251-1255). Verified against gh: issue #340 OPEN ('re-runs the full (paid) segmenter anyway'), plus #330 (MD5 segmentation reuse) OPEN.

65 - why it matters: A page refresh or SPA 'Try again' on an already-indexed video (same MD5) re-bills Gemini (~\$0.10/8-min match) and burns a multi-minute compute run for zero new information. Live-triggered once in the alpha (2026-06-22, per #340).

66 - proposed fix: In _execute_processing, when db.get_video(video_id) exists with status='processed' and existing play_segments (and the request doesn't explicitly force reprocess), return the stored segments as the job result instead of dispatching the segmenter; add a `force` flag to ProcessRequest for deliberate re-ingest.

67 - verifier notes: Severity P1 is honest: concrete failure scenario is a user-reachable one-click ("Try again" after the SPA's 10-min poll ceiling, live-triggered 2026-06-22 per #340) that silently re-bills a metered Gemini key ? and per #340 the M8 cost ledger does not meter Gemini yet, so the re-bill is invisible. Proposed fix is sane and correctly scoped: the skip condition must require status=="processed" AND existing play_segments (a prior "failed" record ? main.py:1196-1201 ? must still re-run), and the added `force` flag preserves the existing destroy-after-confirm re-ingest path (main.py:1233-1235, 1278). Not owner-gated: the fix reduces real cloud spend rather than incurring it, and the owner already specified the desired behavior + acceptance criteria in OPEN issue #340. Effort M is right (endpoint logic + ProcessRequest field + a no-Gemini-call regression test). One duplicate-tracking note: #330 covers the same reuse mechanism as an optimization; a fix should close both.

68 - verifier reasoning: Verified against origin/master at HEAD e410136 (pulled, "Already up to date"). (1) Code behaves as claimed: backend/main.py:1171-1172 computes `video\_id` = compute_video_id(local_video) then `was\_reingest` = db.get_video(video_id) is not None; the ONLY other use of `was\_reingest` in _execute_processing is main.py:1327, where it is passed to `emit\_indexer\_report` as a report field. There is no reuse/skip/short-circuit branch: after validation passes (main.py:1210-1214 upserts status "processed"), the flow unconditionally dispatches segmentation ? cloud via `\_maybe\_dispatch\_remote` (main.py:1241-1248) or in-process

`segmenter.process\_video` (main.py:1251-1255). ProcessRequest (main.py:313-324) has no `force`/`reprocess` field (only video_path, player_pool, max_frames, segmenter, locale, compute). (2) Not fixed: `git log --oneline -15 -- backend/main.py` shows only unrelated commits (#413, #412, #409, #408, #407, #406, #405, #400, ...); no idempotency/reuse change. (3) Cost claim is real in served config: config.json:26 sets `"default\_segmenter": "gemini"` (the paid path is the production default; the "motion" fallback at main.py:1181-1183 is only the code-level default). (4) Tracking verified via gh: issue #340 OPEN (title and body match the finding verbatim, including the 2026-06-22 live trigger and ~\$0.10/8-min cost) and #330 OPEN (MD5 segmentation reuse).

69

70 ### R1-03 [P1/M] critical-fix ? badminton-highlight-indexer

71 **/api/source re-downloads the ENTIRE remote video from the bucket on every HTTP request (each browser Range request) into a fresh, never-deleted temp dir**

72 - anchor: `backend/main.py:646` · tracked in: ? · finder: engine:correctness

73 - evidence: backend/main.py:644-651: get_source calls

storage.acquire_local_input(rec['filepath'], ...) on EVERY request; storage/__init__.py:89-90: for a non-local ref this does `tempfile.mkdtemp(prefix='rally-input-')` + full fetch ? a new dir + full download per call, never removed. Proxy warming is explicitly local-only (main.py:655-657), so a remote-source video never gets a proxy. FileResponse under fastapi>=0.111 (pyproject.toml:28) serves Range ? a browser scrubbing a <video> issues many separate requests to the same URL.

74 - why it matters: Direct cloud-egress cost leak plus a disk-fill denial-of-service of the appliance, reachable through a supported flow (remote-registered assets + P8 annotate-in-UI).

75 - proposed fix: Cache the fetched local copy keyed by video_id (e.g. under <output>/proxies or a per-video workspace source/ dir) and reuse it across requests; extend proxy warming to remote sources; delete or reuse the scratch dir instead of mkdtemp-per-request.

76 - verifier notes: Path correction: the storage citation is backend/storage/__init__.py:89-90 (not storage/__init__.py). Sharper failure scenario: store a 4 GB match via POST /api/assets with an s3://gs:// medium (P4, #409), open the P8 annotate UI, scrub ~10 times ? ~40 GB bucket egress (~\$4-5 real cost) and ~40 GB of never-deleted rally-input-* temp dirs per session; each seek also stalls behind a full multi-GB download, so the feature is effectively unusable for remote sources. Additional color for the fix: the engine free-space guard (require_free_bytes, #407) is not applied to this download path, and the cached copy should live under a guarded root; the finding's proposed fix (cache keyed by video_id, reuse across requests, extend proxy warming to remote) is sane and correctly scoped.

77 - verifier reasoning: backend/main.py:646 calls storage.acquire_local_input(rec["filepath"], ...) on every GET /api/source/{video_id} when no proxy exists (proxy check main.py:641-642). backend/storage/__init__.py:89-90: non-local refs get a fresh tempfile.mkdtemp(prefix="rally-input-") + full fetch per call; docstring (lines 80-82) admits the copy "relies on OS temp reclamation" (nonexistent on Windows %TEMP%), and a repo-wide grep for "rally-input" shows no cleanup path. Remote filepaths are a supported state: POST /api/assets registers the cloud stored_uri as filepath (main.py:838, 855; comment at main.py:1680 lists gs://, s3://, gdrive://). Proxy warming is explicitly local-only (main.py:655-657, docstring 628-629 "remote-source proxying is a follow-up"), so remote sources never escape the full-fetch path. Installed fastapi 0.136.1 / starlette 1.0.0 FileResponse honors Range, so a scrubbing <video> issues many independent requests, each a full-object download. git log -15 -- backend/main.py (through #413) shows no fix after #406 introduced this. Not owner-gated: the fix is plain engineering that reduces cloud spend.

78

79 ### R1-04 [P1/M] tech-debt ? badminton-highlight-indexer

80 **Import cycles: extracted API modules reach back into backend.main globals at call time; config layer imports pipeline layer**

81 - anchor: `backend/api/job\_worker.py:75` · tracked in: docs/CODE_CLEANUP_AUDIT_2026-06-24.md P5 note + D1/P10 follow-up (line 221) · finder: engine:structure

82 - evidence: AST scan of backend/ found 5 strongly-connected components. (1) backend.main <-> backend.api.job_worker <-> backend.api.uploads: job_worker.py:75,98,214,234,251 and uploads.py:47 do call-time 'import backend.main as _main' to read main.py's module globals db_instance/global_config (main.py:279-280); main.py:292 re-exports job_worker names back onto itself. Documented as a transitional shim (CODE_CLEANUP_AUDIT P5 note, line 221) but it makes main.py load-bearing for every extracted module and keeps backend.main AND backend.worker.__main__ mypy-quarantined (mypy.ini:15-23). (2) Layering inversion: backend/config/models.py:249 imports segmenter_registry from backend.pipeline.segmenters.base inside a validator, while base.py:15 (_normalize_config) imports AppConfig back from

config.models ? config validation now depends on pipeline registration side effects. (3-5) eval-package cycles (golden_fixtures<->golden_real_fixtures, ablation<->ablation_pdf, annotations<->gt_loader) are lazy and explicitly documented (golden_fixtures.py:1005-1028) ? lower priority.

83 - why it matters:

84 - proposed fix: Kill the shim as the concrete step-2 of D1: move db_instance/global_config onto app.state (create_app already exists per P5) and pass config/db explicitly to job_worker/uploads, then lift backend.main and backend.worker.__main__ out of mypy quarantine. For the layering inversion, move the segmenter-name validation out of the pydantic validator into a startup check (backend/config/startup.py) so config/ no longer imports pipeline/.

85 - verifier notes: Mechanism correction to the evidence: job_worker's five call-time imports do NOT read db_instance/global_config ? job_worker already takes config/db as explicit params (job_worker.py:89,205,230,247; mypy.ini:14 notes it was lifted from quarantine) and resolves FUNCTION seams (ProcessRequest, _execute_compile, _execute_processing, JobWaiting, _arm_wake_timer, _get_executor, _MAX_DRAIN_WAKE_SEC) through the backend.main module object per the documented test-monkeypatch contract (job_worker.py:8-16, main.py:286-291). Only uploads.py:49 reads a main global (global_config). So the app.state migration for db/config is mostly about uploads + main.py itself; the harder remaining work is redesigning the monkeypatch seam contract and its tests ? effort is top-end M. Sharper failure scenario: any alternative entrypoint (new console_script, uvicorn factory) that builds the app without running main() leaves backend.main.global_config=None and uploads silently falls back to default upload dir/size limits (uploads.py:49-52), ignoring operator security config ? same class as the already-documented python -m backend.main incident (main.py:2188-2195) that hung every queued job. Layering inversion is slightly worse than claimed: the validator's lazy import (models.py:249) transitively executes segmenters/__init__.py:1-7, importing ALL concrete segmenters, so constructing AppConfig anywhere drags the full pipeline in; and SegmenterRegistry.register (base.py:66-70) re-enters AppConfig construction during registration with errors swallowed by base.py:19's bare except.

86 - verifier reasoning: Verified in current tree (HEAD e410136, pulled clean). Cycle 1: backend/main.py:292-302 imports backend.api.job_worker at module load (re-exports); job_worker.py:75,98,214,234,251 do call-time 'import backend.main as _main' (grep-confirmed all 5); uploads.py:47-49 call-time imports and reads _main.global_config; globals defined main.py:279-280, set at main.py:101-103 and 1981-1982; mypy.ini:15-16 quarantines backend.main, :22-23 backend.worker.__main__. Cycle 2: config/models.py:244-256 validator imports segmenter_registry from pipeline/segmenters/base; base.py:14-18 imports AppConfig back in _normalize_config; importing base initializes segmenters/__init__.py:1-7 which imports every concrete segmenter (motion.py:243 etc. register as import side effects), so AppConfig construction depends on pipeline import health ? and register() (base.py:66-70) re-enters AppConfig via segmenter_cls(None, {}) with failures swallowed by base.py:19 bare except. Eval cycles lazy+documented (golden_fixtures.py:1008-1013). Not already fixed: git log for the four files shows only feature commits (#405-#413). Severity honest: the failure class already bit ? main.py:2188-2195 documents the python -m backend.main dual-module incident (backend.main.db_instance=None, drain crashed silently, jobs hung 'queued' forever), patched only for that entry path (main.py:2196-2198); uploads.py:49-52 getattr-defaults mean any entrypoint that skips main() silently runs uploads with default dir/limits. Tracked_in accurate: audit doc line 221 is verbatim the P5 shim note. Not owner-gated: only owner question is D1 phase sequencing, none of the listed gates. Fix sane: create_app exists (main.py:2184), handlers already use app.state (mypy.ini:12-13), backend/config/startup.py exists as the startup-check home.

87

88 ### R1-05 [P1/M] tech-debt ? badminton-highlight-indexer

89 **Path-jail rejects gs://,s3://,gdrive-api:// before scheme resolution ? bucket ingest impossible on the exposed alpha box (issue #327 still open)**

90 - anchor: `backend/main.py:1376` · tracked in: <https://github.com/Khelsutra/badminton-highlight-indexer/issues/327> · finder: engine:security-config

91 - evidence: main.py:1376 calls `validate\_input\_path(req.video\_path, allowed\_root=allowed\_root)` BEFORE `storage.resolve\_input\_ref` at main.py:1379; backend/utils/paths.py:45-55 `resolve\_under\_root` realpath the raw URI so `gs://bucket/x.mp4` escapes the jail -> 400. Startup check startup.py:106-119 makes the jail mandatory when exposed (EXPOSED_NO_VIDEO_ROOT), so an exposed engine can never ingest from a bucket. gh: issue #327 OPEN (verified 2026-07-02), proposing scheme-aware validation.

92 - why it matters: Blocks the alpha CUJ (large-video ingest via gs:// instead of 500 MB browser uploads) and is a stale-open thread since before 06-26; the fix must be done carefully to avoid reopening the arbitrary-local-file hole the jail exists for.

93 - proposed fix: Implement #327 option (b): jail only after resolve_input_ref, applying resolve_under_root exclusively to backend=='local' refs; keep traversal/symlink/absolute rejection for local paths and add both allow (bucket URI accepted) and reject (local traversal still 400) tests, updating tests/test_api_endpoints.py::test_process_hosted_mode_rejects_nonlocal_uri.

94 - verifier notes: Minor corrections: (1) startup check path is backend/config/startup.py:106-119, not startup.py (lines exact). (2) The "500 MB browser uploads" figure in why_it_matters comes from issue #327's body; the repo's actual browser-upload cap is security.max_upload_mb default 4096 MB (backend/api/uploads.py:51, backend/config/models.py:339) ? 500 MB is the Gemini provider per-clip cap (backend/providers/gemini.py:23). Immaterial to the finding (multi-hour matches exceed 4 GB anyway; docs/VIDEO_LOCALITY_MODEL.md:26 "The 10 GB problem is real"). (3) Fix scope additions: there is a second resolve call-site (CLI --video-path, see tests/test_api_endpoints.py:922) ? verify whether it needs the same scheme-aware treatment; ensure unknown schemes still ValueError?400 via resolve_input_ref as the sole scheme gate. (4) Caution, not a gate: the change relaxes a deliberately pinned security contract on the exposed box; #327 records the owner's explicit want, but the option (a) vs (b) choice and the PR should get owner sign-off before landing since it alters exposed-serving posture.

95 - verifier reasoning: Personally reproduced from current code. (1) Ordering: backend/main.py:1376 calls validate_input_path(req.video_path, allowed_root=allowed_root) BEFORE storage.resolve_input_ref at main.py:1379-1381; ValueError maps to HTTP 400 at main.py:1382-1383. (2) Mechanism: backend/utils/paths.py:58-68 routes any request with allowed_root set through resolve_under_root, which os.path.realpath()s the RAW string (paths.py:47) ? 'gs://bucket/x.mp4' realpath to a cwd-relative path, commonpath != root (paths.py:50-54) ? ValueError ? 400. The code itself documents this at main.py:1387-1389: "a configured allowed_video_root (hosted mode) rejects a non-local URI as out-of-root". (3) Pinned by test: tests/test_api_endpoints.py:903-912 asserts gs://, s3://, gdrive:// all ? 400 when allowed_video_root is set ("Pins the security contract"). (4) Exposure mandate: backend/config/startup.py:106-119 emits LEVEL_ERROR EXPOSED_NO_VIDEO_ROOT when edge_auth_enabled and allowed_video_root unset ? so an exposed engine must have the jail, making bucket ingest impossible there. (5) Not fixed: git log -15 -- backend/main.py (head e410136, #413) shows no commit addressing this; behavior present in current tree. (6) gh issue view 327: state OPEN, created 2026-06-22, title/body match the finding exactly, including the owner's explicit want ("ingest large videos via gs:// on the hosted box") and options (a)/(b). Concrete failure: exposed alpha box (edge_auth_enabled=true, allowed_video_root=/jail per boot requirement) receives POST /api/process {"video_path":"gs://bucket/match.mp4"} ? 400 "path ... escapes the allowed root"; only ingest path left is the capped browser upload (backend/api/uploads.py:51, security.max_upload_mb) ? blocks the tracked alpha large-video CUJ. Severity P1 honest (owner-wanted CUJ blocked, open since 2026-06-22, workaround capped/painful per docs/VIDEO_LOCALITY_MODEL.md:19-26). Proposed fix (option b: jail only backend=='local' after resolve_input_ref, keep null-byte check universal and traversal/symlink jail for local, add allow+reject tests, update the pinned test) matches the issue's own proposal and is correctly scoped; effort M is right.

96

97 ### R1-06 [P1/S] doc-rot ? khelsutra (avidullu/khelsutra)

98 **Docs claim the retired droplet alpha is LIVE (droplet deleted 2026-06-25)**

99 - anchor: `NEXT\_STEPS.md:12` · tracked in: partially in issue #97 (item 2: update DEPLOY_ALPHA.md); the NEXT_STEPS.md droplet claim is untracked · finder: khelsutra:issues-triage

100 - evidence: NEXT_STEPS.md:12-13 ('Where we are (updated 2026-07-01)') still says the site is 'live + browsable on the D0 droplet -> http://168.144.126.176:8787'; deploy/DEPLOY_ALPHA.md:3 says 'Status: LIVE (deploy live-verified on the droplet 2026-06-22)' and :29 'This runbook ships the split topology'; docs/LOCAL_APPLIANCE_ARCHITECTURE.md:176 'The alpha stand-up is now LIVE (2026-06-22)'. Project memory (MEMORY.md:63, session-handoff.md:19): droplet claude-do-rally-test/168.144.126.176 RETIRED+deleted 2026-06-25; alpha paused pending Cloud Run. nslookup shows api./app.khelsutra.guru still resolve via Cloudflare to a tunnel with no origin (pending CNAME/Access cleanup).

101 - why it matters:

102 - proposed fix: One doc PR: mark DEPLOY_ALPHA.md status RETIRED/PAUSED with a pointer to the Cloud Run plan, fix NEXT_STEPS.md 'Where we are' (droplet gone, site is CF-Workers-only), and correct LOCAL_APPLIANCE_ARCHITECTURE.md:176. Also note the single-origin repoint that superseded the split topology (issue #97). The 2026-07-01 refresh (#139) touched docs but kept the stale claim, so this will keep misleading ramp-up reads.

103 - verifier notes: Minor anchor correction: the 'Where we are (updated 2026-07-01)' header is NEXT_STEPS.md:5; the droplet-live claim spans lines 12-13 (the finding's line 12 anchor is correct for the URL text). Severity P1 stands: NEXT_STEPS.md's 2026-07-01 stamp is

NEWER than the handoff's 2026-06-25 correction, so ramp-up readers reconciling the two will trust the wrong doc; worse, DEPLOY_ALPHA.md \$1 (line 35) directs deploying the engine + secrets onto 168.144.126.176 ? an IP the owner no longer controls and DO can recycle. Fix is sane, doc-only, and not owner-gated (it records already-made owner decisions: retirement 2026-06-25 + single-origin repoint per issue #97); one caution ? the fix should point to the Cloud Run move as 'paused pending' rather than declare a new serving decision, which would be owner-gated. Did not independently re-verify the finder's nslookup/DNS claim; the finding stands without it.

104 - verifier reasoning: Working tree at HEAD f9eea14: NEXT_STEPS.md:5 'Where we are (updated 2026-07-01)' + :12-13 'live + browsable on the DO droplet -> http://168.144.126.176:8787 (systemd khelsutra-site)'; deploy/DEPLOY_ALPHA.md:3 'Status: LIVE (deploy live-verified on the droplet 2026-06-22...)' + :29 'This runbook ships the split topology'; docs/LOCAL_APPLIANCE_ARCHITECTURE.md:176 'The alpha stand-up is now LIVE (2026-06-22): the engine booted on the droplet (168.144.126.176)'. Ground truth: memory/khelsutra-site-droplet-parity.md:3,10-14 (droplet DELETED 2026-06-25, data archived, doctl delete done, monitor removed via PR #117), session-handoff.md:19 ('the prior LIVE alpha on the droplet is GONE ? alpha paused pending the Cloud-Run move'), MEMORY.md:63. Not fixed: git log shows the last commit touching all three files is 9223ba9 (#139, the 2026-07-01 refresh) and the stale text persists at HEAD. Tracking verified: gh issue 97 is OPEN, item 2 covers only DEPLOY_ALPHA.md (single-origin update); the NEXT_STEPS.md droplet claim is untracked. Issue #97 also confirms split topology was superseded by the 2026-06-22 single-origin repoint, so DEPLOY_ALPHA.md:29 is doubly stale.

105

106 ### R1-07 [P1/M] critical-fix ? khelsutra (avidullu/khelsutra)

107 **Issue #111 still fully valid: refresh loses the in-flight job; progress static for minutes**

108 - anchor: `web/src/hooks/useIngestJob.ts:117` · tracked in: khelsutra#111 · finder: khelsutra:issues-triage

109 - evidence: useIngestJob.ts:117-128 ? job_id lives only in the poll closure (no localStorage/URL persistence); App.tsx:71-73 seeds state only from ?video= (set on done, App.tsx:153). A refresh mid-job lands the user back on the ingest form while the server job keeps running. Progress is the coarse free-text job.progress (useIngestJob.ts:106); no heartbeat/animation beyond a spinner chip (IngestPanel.tsx:337-358).

110 - why it matters:

111 - proposed fix: SPA half is Claude-doable now: persist the active job_id (?job=<id> + localStorage) and on load resume polling to the existing ?video= handoff; add elapsed-time/heartbeat copy. Engine half (granular 'segmenting chunk 3/7' from backend/pipeline/segmenters/gemini.py via the job-progress hook) is a small cross-repo change. Per the issue this is 'the most likely way an early tester loses a working job'.

112 - verifier notes: Failure scenario, sharpened from issue #111's field report: an ~8-min upload spends ~5 min in Gemini segmentation with near-static "processing" text; the tester refreshes (natural when nothing moves), lands on the ingest form with no trace of the running job, and either re-submits ? a duplicate job, potentially on paid cloud GPU via the compute picker ? or abandons believing it failed; the finished video only appears later in the library with no pointer back. Fix note: the backend is already refresh-safe (jobs persisted in SQLite), so the SPA half (persist ?job=/localStorage, resume polling into the existing ?video= handoff, heartbeat/elapsed-time copy) is standalone and Claude-doable now; only the granular-progress half touches the sibling badminton-highlight-indexer repo.

113 - verifier reasoning: useIngestJob.ts:117-128 ? job_id is a closure-local (`let jobId` at :117, used only by poll(:128)); grep of web/src confirms no localStorage/URL persistence of any job id (only annotate flag, wizard key, i18n lang). App.tsx:71-73 seeds videoId solely from ?video=, written only on job-done via loadVideo (App.tsx:150-154 ? useIngestJob.ts:96-98), so a mid-job refresh renders the ingest form (App.tsx:338-343) while the server job keeps running. Progress is coarse free-text job.progress (useIngestJob.ts:106) shown as a spinner chip (IngestPanel.tsx:337-348), no heartbeat/elapsed time. git log -15 on both files (latest #140, #138, #128?) shows no fix; issue khelsutra#111 is OPEN and documents the identical gaps from a live 2026-06-22 alpha run. Engine path backend/pipeline/segmenters/gemini.py exists in the sibling badminton-highlight-indexer repo as the finder scoped it. Not owner-gated: pure client resiliency, no go-live/corpus/licensing/spend/legal decision.

114

5. user (2026-07-02T04:45:57.032Z)

190 - evidence: \$7 tracker (lines 63-73): P1-P8 all ?. Verified in code: grep 'build_stitch_plan' across backend/ and tests/ returns nothing (P2's pure stitch-plan seam was never factored out of backend/pipeline/stitcher/compiler.py), and tests/test_golden_fixtures.py

has no stitch-plan snapshot. Minor rot: the P0 row (line 65) still shows ? '(this PR)' though the doc merged as #383 (commit 1283e18, 2026-06-25) and the header still says 'DRAFT'.

191 - why it matters: This gate exists precisely to bless genuine (non-AST-identical) refactors as quality-neutral ? i.e., exactly the class of change as the stalled #390 main.py split. Executing P1-P4 first would de-risk redoing the god-module split; leaving it a week-stale locked-decisions doc means refactors continue with no behavioral proof.

192 - proposed fix: Start \$7 P1+P2 (canonical decision-snapshot schema + factor pure build_stitch_plan() out of compiler.py) as the next Claude-doable engine work; flip the P0 row to ? #383 and status to IN PROGRESS in the same PR. Ideally land P1-P4 before re-cutting the main.py split.

193 - verifier notes: One factual correction to the evidence: P8 (line 73) is ? (blocked/owner-gated), not ? ? "P1-P8 all ?" should read "P1-P7 all ?, P8 ?". Substance unaffected: zero of nine tracker rows show any progress. Sharper failure scenario: #390 gets re-cut and merged as a "quality-neutral" god-module split with no behavioral proof ? the only offline gate today is the segmentation-only characterization snapshot (test_replay_matches_committed_expected), so a subtle behavior change in stitch/concat-plan construction (clip order, cut points, padding) ships undetected until a GPU nightly or a user-visible reel regression. Effort L is honest for the proposed P1-P4 scope; the minimal P0-row/status flip alone is S and could be split out.

194 - verifier reasoning: Personally verified in the current working tree (master, clean). (1) docs/DECISION_SNAPSHOT_REGRESSION.md:8 still reads "Status: `DRAFT` ... LOCKED (2026-06-25) ... Last updated: 2026-06-25"; line 65 P0 row is "? | (this PR)" even though the doc merged as PR #383 (gh: MERGED 2026-06-25T17:47:44Z; sole commit 1283e18 touching the file ? `git log --oneline -15 -- docs/DECISION\_SNAPSHOT\_REGRESSION.md` returns exactly one commit, so nothing fixed it since). (2) Zero code shipped: grep for build_stitch_plan across the repo hits only the doc itself and docs/README.md; backend/pipeline/stitcher/compiler.py:24 contains only `class VideoCompiler` with no pure plan-builder seam; grep -i stitch under tests/ returns 11 files NOT including tests/test_golden_fixtures.py, confirming no stitch-plan snapshot exists; repo-wide grep for stitch_plan|decision_snapshot matches only docs. Tracker rows at lines 66-72 (P1-P7) are all ? with PR "?". (3) Why-it-matters context holds: gh shows PR #390 ("refactor(D1/P10): extract orchestration.py from main.py (god-module split, step 1)") is still OPEN ? exactly the non-AST-identical refactor class this gate (D11, line 40) is designed to bless. (4) Not owner-gated: the only owner-gated rows are P0 (already satisfied ? line 82 "[x] D1?D11 locked by the owner (2026-06-25)") and P8 (line 73); P1-P7 are marked "Gated? No", and the proposed fix touches only non-gated rows and none of the listed gate categories.

195

196 ### R1-17 [P2/M] tech-debt ? badminton-highlight-indexer

197 **Fetched-input scratch copies are never reclaimed anywhere (process jobs, compile jobs, /api/assets, cloud_run marker reads) ? 'OS temp reclamation' premise is false on Windows**

198 - anchor: `backend/storage/\_\_init\_\_.py:89` · tracked in: ? · finder: engine:correctness

199 - evidence: storage/__init__.py:80-90: acquire_local_input docstring says the scratch copy 'is left under the system temp dir ... and relies on OS temp reclamation'. Call sites that leak one full-video copy per run: _execute_processing (main.py:1168-1171), _execute_compile (main.py:1683-1686), create_asset (main.py:791-799), get_source (main.py:646). cloud_run.py:219-222 additionally leaks a cr-read-* mkdtemp per marker/report read. No janitor exists (grep for rally-input-/cleanup found none); engine issue #301 covers only the upload_dir retention sweep, not temp scratch.

200 - why it matters: Slow-burn disk exhaustion on exactly the appliance the product ships as; also inflates the free-space guard's false rejections.

201 - proposed fix: Wrap remote-input acquisition in a per-job scope that deletes the scratch dir in a finally (worker already owns the job lifecycle), or route scratch under output_base(cfg)/tmp with an age-based sweep at startup; make cloud_run._read_json use TemporaryDirectory.

202 - verifier notes: Two sharpenings. (a) get_source (main.py:646) cannot use a simple finally-delete: the FileResponse streams from the scratch file after the handler returns, so deletion there needs the age-based-sweep option (or a keyed scratch cache reused across views) ? the finder's second proposed option covers it; per-job finally scopes fit _execute_processing/_execute_compile/create_asset. (b) Add a fifth leaking call site to the fix scope: the CLI single-shot path at main.py:2056. Sharper failure scenario: with input_backend=gdrive-api/gcs, each processed match leaks one full multi-GB copy into %TEMP% per process job, another per compile, one per /api/assets source_uri registration (fetched solely to hash), and one per /api/videos/{id}/source view ? until require_free_bytes(tempfile.gettempdir()) at main.py:789/1402 starts rejecting all further remote ingests at stage="fetch" even though the disk is full of the tool's own dead scratch.

203 - verifier reasoning: Personally verified in the current tree. (1)

backend/storage/__init__.py:80-82 docstring says the scratch copy "is left under the system temp dir ? and relies on OS temp reclamation"; line 89 does tempfile.mkdtemp(prefix="rally-input-") and line 90 fetches into it ? no deletion. (2) Every claimed call site leaks: _execute_processing main.py:1168-1171 (no rmtree anywhere in main.py ? grep confirmed rmtree exists only in eval/golden_real_fixtures.py, pipeline/stitcher/compiler.py, pipeline/detectors/*); _execute_compile main.py:1683-1686; create_asset main.py:791-799 fetches a full video copy just to MD5-hash it (main.py:805) then returns at 865-875 without deleting local_working; get_source main.py:646 fetches per-request and serves it via FileResponse (line 660) ? a per-view full-video leak for remote sources. Plus a fifth site the finder missed: CLI single-shot main.py:2056. cloud_run.py:219-222 _read_json mkdtemps "cr-read-" per marker/report read, no cleanup. (3) No janitor: tools/cleanup_caches.py scans only WSL ~/clips, the output/tree, and security.upload_dir (docstring lines 3-5, classify() lines 52-97) ? never tempfile.gettempdir(); its "tmp*" rule (line 79-80) applies only inside those trees. (4) Not fixed: git log for storage/__init__.py tops out at b6214e5 (repo-wide ruff format); gh issue 301 confirmed to cover only the T11 upload-retention scheduled job, not temp scratch. (5) Guard interplay real: require_free_bytes(tempfile.gettempdir(), ?, stage="fetch") at main.py:789 and main.py:1402 ? leaked copies shrink exactly the headroom the fetch guard checks, so accumulation converts into false fetch rejections. Windows does not auto-reclaim %TEMP% by default, so the docstring premise is false on the shipping platform. Severity P2 honest: affects only remote-input flows (local paths are identity passthrough, __init__.py:85-88), slow-burn, no data loss. Not owner-gated ? pure disk hygiene, no strategy/spend/ToS/corpus decision.

204

205 ### R1-18 [P2/S] critical-fix ? badminton-highlight-indexer

206 **GDriveApiStorage.fetch_to_local buffers the entire file in RAM
(get_media().execute()); put() deletes the old object before the new upload succeeds**

207 - anchor: `backend/storage/gdrive\_api.py:144` · tracked in: ? · finder:
engine:correctness

208 - evidence: gdrive_api.py:143-148: `data = self.\_service.files().get\_media(fileId=file\_id).execute()` then `fh.write(data)` ? the whole object is materialized as one bytes value (no MediaIoBaseDownload streaming).
gdrive_api.py:163-170: put() resolves an existing file at the path and `files().delete(...)` BEFORE `files().create(...)` ? delete-then-create is not atomic.

209 - why it matters: Latent OOM + data-loss on the Drive-API storage hot path before it goes live for the reel round-trip.

210 - proposed fix: Stream downloads with googleapiclient.http.MediaIoBaseDownload in chunks; for put(), create the new file first (or upload to a temp name) and delete the old one only after success.

211 - verifier notes: File/line accurate (line 144 is the get_media().execute() call).
Failure scenarios: (1) fetch_to_local on a multi-GB source/reel buffers the full file in RAM on a memory-limited Cloud Run box ? OOM kill; (2) re-put() of a reel to an existing path deletes the old file, then the resumable create fails (network/quota) ? the reel is gone with no copy at either end. Fix refinement: for put(), prefer `files().update(fileId=existing, media\_body=...)` when a file already exists ? atomic in-place content replacement that preserves the file id and avoids both the delete-first data-loss window and the duplicate-(name,parent) window that create-first-then-delete would open (the file's own comment at lines 160-162 notes _resolve_id returns the first match among duplicates). Download fix as proposed: MediaIoBaseDownload chunked streaming.

212 - verifier reasoning: backend/storage/gdrive_api.py:143-148 ? fetch_to_local does `data = self.\_service.files().get\_media(fileId=file\_id).execute()` then `fh.write(data)`: the whole object is returned as a single bytes value; no MediaIoBaseDownload/chunked loop exists in the file (only MediaFileUpload for uploads, lines 78-84). Docstring (lines 3-6, 19) states this backend serves videos/reels on headless Cloud Run/GPU hosts, so multi-GB buffering is a genuine OOM risk. backend/storage/gdrive_api.py:163-170 ? put() resolves `existing`, calls `files().delete(fileId=existing).execute()`, and only then `files().create(...)` : a create failure after the delete loses the prior object (non-atomic overwrite). git log -- backend/storage/gdrive_api.py shows only formatting/isort commits (b6214e5, f523d3a) since the introducing commit 819e395 (M12, PR #338) ? not fixed. Not owner-gated: the module's owner gate (lines 14-17) covers the per-user OAuth flow, not this code's download/upload mechanics.

213

214 ### R1-19 [P2/S] critical-fix ? badminton-highlight-indexer

215 **config.json output.output_dir is silently clobbered by the CLI --output-dir DEFAULT on every server/CLI start ? the typed knob is dead, and the DB path moves with it**

216 - anchor: `backend/main.py:1992` · tracked in: ? · finder: engine:correctness

217 - evidence: main.py:1837-1841 defines `--output-dir` with default='output'; main.py:1992 unconditionally does `global\_config.output.output\_dir = resolve\_output\_dir(args.output\_dir)` ?

the override fires even when the flag was never passed, so a config.json value can never take effect. models.py:274-276 documents the knob as if it's honored ('The CLI's --output-dir overrides this at startup' ? implying only when given). db_path derives from it (main.py:1995-1997).

218 - why it matters: A documented config knob that is a no-op on the main entry path, with data-location consequences (perceived data loss / split state).

219 - proposed fix: Change argparse default to None and only override when args.output_dir is not None (falling back to the config value through resolve_output_dir).

220 - verifier notes: Path correction: the documented knob is at backend/config/models.py:274-276, not backend/models.py (which does not exist). Sharper failure scenario: user sets output.output_dir: "D:/rally-data" in config.json and restarts `indexer --app`; main.py:1992 resets the dir to resolve_output_dir("output"), a fresh empty DB is created there (main.py:1995-1997) so the UI shows zero videos (perceived data loss), while run_eval.py/run_regression.py via backend/config/__init__.py:41-42 resolve the DB at D:/rally-data ? server and eval tooling disagree on DB location. Fix scope addition: the trim path at backend/main.py:1927-1929 also calls resolve_output_dir(args.output_dir) before config load and must handle a None default (fall back to the loaded config value there too, or to "output" since global_config isn't initialized yet at that point). Effort stays S.

221 - verifier reasoning: Read current tree. backend/main.py:1837-1841 defines --output-dir with default="output". backend/main.py:1982 loads the typed config; backend/main.py:1992 then unconditionally does global_config.output.output_dir = resolve_output_dir(args.output_dir), which fires even when the flag was never passed. resolve_output_dir (backend/utils/app_home.py:83-100) never consults the config value ? with RALLY_APP_HOME unset it returns the argparse default "output" verbatim ? so a config.json output.output_dir can never take effect. db_path derives from the clobbered value (backend/main.py:1995-1997). The only production entry is main() ? create_app (main.py:2184) ? uvicorn.run (main.py:2185); there is no module-level app, so every server/CLI start hits the clobber. The knob is documented at backend/config/models.py:274-276. Not fixed: git log --oneline -15 -- backend/main.py tops out at e410136 (#413), none touching this. Split-state is concrete: backend/config/__init__.py:41-42 (default_db_path, used by run_eval.py/run_regression.py) builds the DB path directly from config.json's output.output_dir with no clobber ? so with the knob set, the server and the eval tooling resolve DIFFERENT DB locations. Not owner-gated (local config/CLI startup only).

222

223 ### R1-20 [P2/S] critical-fix ? badminton-highlight-indexer

224 **upload_dir default 'output/uploads' is CWD-relative and never app-home-resolved ? packaged --app writes uploads into an arbitrary CWD (or 500s on an unwritable one) while DB/reels live under RALLY_APP_HOME**

225 - anchor: `backend/api/uploads.py:50` · tracked in: ? · finder: engine:correctness

226 - evidence: uploads.py:50 `upload\_dir = getattr(sec, 'upload\_dir', None) or 'output/uploads'` used verbatim at :89 (os.makedirs), :95, :170 ? grep confirms no resolve_output_dir/app_path is ever applied to it (only output.output_dir is resolved, main.py:1992 via app_home.py:83-100). models.py:336 default 'output/uploads'.

227 - why it matters: Breaks or misplaces the primary ingest path of the packaged app; exactly the class of CWD-vs-app-home bug P0a was built to eliminate (app_home.py:1-8).

228 - proposed fix: Resolve upload_dir through resolve_output_dir()/app_path() in _upload_cfg (relative ? app-home-rooted; absolute passes through), mirroring output_dir.

229 - verifier notes: Sharper failure scenario: on macOS the .command runs with Terminal's CWD=\$HOME, and any `indexer --app` shell launch (advertised in install.py:243) uses the caller's CWD ? uploads land in e.g. ~/output/uploads while DB/reels live under ~/Library/Application Support/Khelsutra/output; unwritable CWD ? PermissionError at uploads.py:89 ? 500 on /api/upload/init. Caveat: Windows double-click of Khelsutra.bat coincidentally works (the .bat lives in app_home, install.py:175, and cmd sets CWD to the .bat dir). Extra impact the finder missed: tools/cleanup_caches.py:403 sweeps <output-dir>/uploads (app-home-rooted post-P0a) and the uninstall manifest user_data (install.py:117-121) lists only app-home paths, so stray multi-GB uploads are never swept and survive uninstall --purge. Fix as proposed is sane: resolve_output_dir() in _upload_cfg is byte-identical with RALLY_APP_HOME unset (app_home.py:94-95) and existing tests use absolute tmp paths (test_upload_download.py:31) that pass through.

6. assistant (2026-07-02T04:46:01.270Z)

Now let me read the repo's CLAUDE.md and CODE_MAP.

7. user (2026-07-02T04:46:02.190Z)

```

1      # CLAUDE.md ? project memory / agent handoff
2
3      Local, AI-powered badminton (? racket-sports) highlight indexer + rally detector**:
4      FastAPI + SQLite + ffmpeg + GPU CV/AI, with paid Gemini as the cloud AI. Sibling
5      repos: `sports-data-collector`, a VLC-based `rally-annotator`.
6
7      ## Read these first
8      - docs/README.md ? doc index. docs/CODE_MAP.md ? code navigation + data
9        contracts (the cold-start map; read before touching code).
10     - Authoritative owned-model roadmap: docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md (+
the
11         live docs/NEXT_STEPS.md). These SUPERSEDE the 2026-06-07 strategy docs below where
they
12         disagree (licensing, Gemma-4 stance, base model, conversational-QA).
13     - Strategy (post-scaffolding): docs/COMMERCIALIZATION.md,
docs/PMF_MARKET_RESEARCH.md,
14         docs/PLATFORM_ARCHITECTURE.md, docs/DEFERRED_HARDENING_PLAN.md. Also
docs/ROADMAP.md
15         (offline-segmenter narrative) and docs/BACKLOG.md.
16     - Rally-detection quality track (design merged 2026-06-13, owner-gated, implementation
in progress):
17         docs/HOW_RALLY_DETECTION_WORKS.md (2-layer mental model) ?
docs/MULTI_SIGNAL_FUSION_PLAN.md
18         (fuse shuttle + person + audio cues; the held-shuttle bug is the already-built
rally_gate.py
19         "Gate A" not wired into production ? cheapest win) ? docs/EXPERIMENT_HARNESS.md
(A/B + shadow
20         eval so cues land flag-gated/default-OFF with zero regression) ?
docs/ROORA_LABELING_GUIDELINES.md
21         (v8 vendor labeling spec). docs/NEXT_STEPS.md has the prioritized pick-up for this
track.
22     - History: docs/archives/ ? ADRs (decisions/DECISIONS.md), research, and
23         checkpoints/ (latest: 2026-06-strategy-foundation/). Archive policy: only
move
24         terminal-state (DONE/REJECTED) work; live docs stay at docs/ top level ? and
before
25         archiving, HONESTLY update the doc first (verify claims vs code ? flip status + a
completion
26         note ? update inbound refs ? keep the lineage ? then git mv), per the checklist +
living
27         lifecycle register in docs/DOC_STATUS.md.
28
29     ## Current state (June 2026)
30     - Scaffolding complete; strategy foundation + owned-model roadmap merged. No
31         platform/owned-model implementation has started. A 2026-06-10 audit-driven hardening
32         sweep landed (licensing gate, detector keying, re-ingest safety, config, eval gate,
API
33         validation, reliability, CI + test coverage) ? see git history.
34     - Two tracks, both owner-gated: the committed next PLATFORM slice = async job
model
35         (DEFERRED_HARDENING_PLAN.md #16, single-tenant); the owned-model track starts at
a
36         greenlit Phase-0 milestone (OWNED_MODEL_IMPLEMENTATION_PLAN.md). Do not start
37         implementation without explicit owner approval.
38
39     ## Architecture in one breath
40     - Pluggable registries (ABC + Registry singleton): segmenters / providers /
validators
41         / sports / annotations ? see backend/pipeline/segmenters/base.py. Future:
42         StorageBackend + ComputeBackend registries (PLATFORM §3).
43     - video_id = MD5 of the file ? backend/utils/hashing.py::compute_video_id (one
helper).
44     - Typed config = backend/config/ (Pydantic). DB schema evolves via the PRAGMA
45         user_version migration ladder in backend/infrastructure/database.py.
46

```

```

47     ## Committed guardrails (non-negotiable)
48     - **Paid LLMs (Gemini/OpenAI/Anthropic) = inference / serving / fallback ONLY ? never a
49       training data source** (terms/copyright). Commercial-clean licensing for EVERY
dependency ?
50     code, data, AND weights: **MIT / Apache-2.0 / BSD only** (ratified 2026-06-09, owner-
agreed).
51     - **Owned-model base:** **Qwen3-VL (Apache-2.0)** primary, InternVL3 (MIT) fallback;
**Gemma-4
52       is AMBER / bench-only** (Apache grant likely but Prohibited-Use posture unconfirmed ?
counsel
53       needed; do NOT build on it without sign-off). **Reject Gemma 3** (viral license).
**Exclude
54       minors** from the training pool; per-user training data needs a separate explicit opt-
in.
55     - **v1 = rally + highlight + history + correction loop**; **gait/biometrics strictly
future**
56       (GDPR Art. 9).
57     - **Generic data schema** ; `coach ? athlete` is an **optional overlay** , not baked in.
58     - **Compute = GCP ONLY (ADR-013, 2026-06-20):** GPU training + serving run on **GCP**
(`deploy/gcp/` ;
59       Cloud Run+GPU for serving per ADR-012). **DigitalOcean is DROPPED for compute** (DO
GPU not enabled
60       on the account + not cleanly credit-funded; Paperspace is subscription-gated). Don't
re-evaluate
61       DO/Paperspace for compute without a new ADR; the `$200` DO credits go to **non-compute
only** . The
62       `deploy/digitalocean/` scripts are kept solely as provider-agnostic Linux-GPU setup
reused by GCP.
63
64     ## Workflow conventions
65     - **Branch off `master` ; ONE PR per work item; NO PR stacking.** Don't open a PR unless
asked.
66     - Tests: `python run_all_tests.py` (offline, fully mocked). Python 3.10+ (the `trackers`
package requires 3.10); `ffmpeg` required.
67     - Never push to the default branch directly; descriptive commit messages.
68     - **Substantial design/project work gets a tracked project doc** following
`docs/PROJECT_DOC_TEMPLATE.md` (status header · decisions-locked · progress tracker ·
definition-of-done; archived on completion). Precedents: `docs/RALLY_QUALITY_RESEARCH.md` §7,
`docs/GOLDEN_REGRESSION_FIXTURES.md` .
69     - **Launch / "bit-flip" PRs** (a default-flip that drives a tracked project to terminal
state) MUST, in the SAME PR, run the `docs/DOC_STATUS.md` §2 archival due-diligence on that
project's docs (verify-vs-code ? flip status + completion note ? update inbound refs ? keep ADR
lineage ? `git mv` to `docs/archives/` ? record in DOC_STATUS §3), and fix any internal-link
breakage the move causes. Enforced by `.github/PULL_REQUEST_TEMPLATE.md` + the `link-check` CI
job (`python -m tools.check_md_links` ).
70     - ?? Hand the heaviest machine-side checks (real video, the full GPU/WSL/WASB suite) to
the owner,
71       who has the golden corpus + labelled videos; verify pure-logic paths locally.
72

```

8. user (2026-07-02T04:46:04.166Z)

On branch master
Your branch is up to date with 'origin/master'.

nothing to commit, working tree clean
e410136 api: /api/videos surfaces fps + resolution (from the stored ffprobe details) (#413)
81a78bf api: /api/videos returns video_id + a friendly match_name (fix disappearing name) (#412)
8b16cf4 docs: reflect today's storage endpoints + the corpus/user-video boundary (#411)

9. user (2026-07-02T04:46:12.828Z)

Already up to date.
e410136 api: /api/videos surfaces fps + resolution (from the stored ffprobe details) (#413)

10. user (2026-07-02T04:46:13.006Z)

No matches found

11. user (2026-07-02T04:46:17.234Z)

```
1      # CODE_MAP.md
2
3      **Purpose:** cold-start navigation map for the badminton-highlight-indexer ? a self-
hosted video ? highlight-reel pipeline (CV/AI rally detection ? SQLite ? ffmpeg stitch) with a
segmenter-agnostic eval/calibration stack, typed config, and per-video observability reporting.
4
5      **Regenerate:** rerun the `build-code-map` workflow (parallel subsystem readers ? this
doc). Re-run after any registry/contract/WSL-interface change, or after a directory restructure
(this revision follows #40).
6
7      ? **2026-06-10 audit-hardening sweep changed several behaviors documented below ? trust
git history over stale lines until the next full regen:** `add_video` is now an **UPSERT** (no
FK-cascade wipe on re-ingest; segmentation is destroy-AFTER-confirm via
`segment_watermarks`/`prune_segments`); detector artifacts are keyed by **video_id**, not
basename; **indexing.wasb.*` is now a real typed config** (was silently dropped); `timeout_sec`
defaults **1800s** (was 0/hang-forever); SQLite opens with **WAL + busy_timeout**;
```

`match\_intervals` requires real overlap (iou>0); `run\_regression` baseline lives in tracked
`eval\_baselines/` and **exits 3 on a missing baseline**;

the API validates
`video\_path`/`output\_filename`. New tests: `test\_api\_endpoints`, `test\_compiler`,
`test\_wasb\_runner`, `test\_reingest\_safety`, `test\_detector\_keying`, `test\_config\_wasb`,
`test\_path\_safety`, `test\_reliability\_hardening` (+ collector `test\_licensing\_gate`,
`test\_import\_guard`, `test\_cvat\_robustness`).

```
8
9      **LEGEND**
10     - `@path` = file (always repo-relative)
11     - `::sym` = symbol (class/func/const) in the nearest preceding `@path`
12     - `->` = dataflow / call / "produces"
13     - `S` = data contract (canonical shape)
14     - `?` = gotcha / footgun
15     - `?` = registry / plugin extension point
16     - `WSL` = runs inside WSL2 conda env, NOT the Windows Python process
17
18     ---
19
20     ## 0. REPO LAYOUT ? where new files go (READ BEFORE committing a new file)
21
22     Top-level tree and the **placement rule** for each kind of file. When unsure, prefer an
existing
23     **registry / extension point** (`?`) over a new top-level file.
24
25     | You're adding? | Put it in? | Notes |
26     |---|---|---|
27     | Backend / runtime code | `backend/<subsystem>/` |
`pipeline/{segmenters,detectors,validators}`, `providers/`, `sports/`, `annotations/`,
`infrastructure/`, `reporting/`, `stitcher/`, `utils/`, `config/`, `features/`, `api/`. Most
additions are **register a class in the relevant `?` registry**, not a new module tree. |
28     | **Eval LIBRARY** (importable, pure ? metrics, calibration, a scorer, a feature) |
`backend/eval/` | Pure core, no side-effects at import; CLI/I-O lives in a separate driver.
These are **imported widely** (e.g. `calibrate_wasb` by 8 sites) ? they are library, not
scripts; do NOT move them out. |
29     | **Standalone run-driver / one-off entry CLI** (has `__main__`, NOT imported by app
code) | **scripts/** | e.g. `scripts/run_phase3_eval.py`, `scripts/run_process_and_stitch.py`.
Run as `python scripts/<name>.py`. If it needs `load_dotenv()`/`sys.path` before importing
backend, add `scripts/<name>.py" = ["E402"]` to `ruff.toml`. |
30     | **Ops / maintenance tool** (not part of the served app ? cache GC, batch admin) |
`tools/` | Run as `python -m tools.<name>` (it's an importable package). e.g.
`tools/cleanup_caches.py`. Keep the destructive decision path pure + unit-tested. |
31     | Training code (model fitting) | `training/` | Behind the CI-enforced import wall
(ADR-008) ? `backend.main` must not import it. Own `requirements-train.txt`. |
```

```

32      | A test | `tests/test_*.py` | pytest auto-discovers; the full-suite lane has a **70%
coverage floor**. Pure/offline + mocked (no GPU/WSL/network). |
33      | A doc (design, plan, living log) | `docs/` | Index it in `docs/README.md`. Only
**terminal-state** (DONE/REJECTED) work moves to `docs/archives/`. |
34      | A throwaway repro / debug script | `scratch/` | **Gitignored** + excluded from
ruff/pytest. Never imported by app code. |
35      | A model artifact (weights + manifest) | `artifacts/<name>/<ver>/` | ADR-008:
**manifest committed** (provenance), **weights gitignored**. |
36      | Eval baseline / leaderboard JSON | `eval_baselines/` | Tracked (survives a clean
checkout); the no-regression gate reads it. |
37      | **Committed regression fixture** (video-free golden) |
`eval_baselines/fixtures/<slug>/` | Tracked, **LF-pinned** (`gitattributes`). Synthetic Tier-0
(or owner-cleared, de-attributed real). Drives the `golden_fixtures` characterization gate with
**no video/GPU/DB**. See `docs/GOLDEN_REGRESSION_FIXTURES.md`; mint via `python -m
backend.eval.golden_fixtures --freeze-synthetic` / `--freeze-real` (refusal-gated). |
38      | Pipeline run outputs (reels, trajectories, DBs, frames) | `output/` | **Gitignored.**
Per-video isolation is the tracked **`PER_VIDEO_WORKSPACE.md`** project (P0 helper shipped #264;
supersedes the archived `PER_VIDEO_OUTPUT_LAYOUT`). |
39
40      **Canonical files that stay at the repo ROOT** (referenced by CI / packaging / docs ? do
not move):
41      `pyproject.toml` (PEP 621 packaging, hatchling; `indexer` entry point) .
`run_all_tests.py` (CI: `ci.yml`) .
42      `run_regression.py` (the no-regression gate CLI, cited in `EXPERIMENT_HARNESS.md`) .
`run_eval.py`
43      (grandfathered ? imported by `tests/test_eval_harness.py`). The old diagnostic
`setup.py` moved to
44      `scripts/doctor.py` (it shadowed the build system); invoke via `python
scripts/doctor.py` or `indexer --setup`.
45      New standalone run-drivers go in **`scripts/`**, not the root.
46
47      ---
48
49      ## 1. PIPELINE
50
51      ### 1A. Runtime: video file -> highlight reel
52      ```
53      typed config load (built-in defaults -> config.json overrides; absent/parse-fail -> all-
defaults, never fatal)
54
55      @backend/config/loader.py::load_config ->
@backend/config/models.py::AppConfig
56      -> video_id = MD5(whole file) @backend/utils/hashing.py::compute_video_id
(single shared helper, 64KB chunks; imported by main + all run_*.py ? no more per-entrypoint
copies)
57      -> validate
@backend/pipeline/validators/base.py::registry.run_all_with_context(video_path, global_config)
58      FormatValidator -> ctx['ffprobe_meta'] -> ResolutionFPSValidator -> PTSContinuity
-> SceneCut -> Blur -> Relevance (? ACTUAL registration/exec order ? see
@backend/pipeline/validators/__init__.py)
59      [validation FAIL -> persist status="failed", return early; report STILL emitted
in finally]
60      -> db.add_video(video_id, filepath, status, [r.dict()])
@backend/infrastructure/database.py
61      -> seg_name = req.segmenter or global_config.indexing.default_segmenter (fallback
'motion')
62      -> seg = segmenter_registry.get(seg_name)(db, global_config) ?
@backend/pipeline/segmenters/base.py::segmenter_registry (400 + available list if missing)
63      -> segments = seg.process_video(video_id, video_path, player_pool?, max_frames?)
[STOP POINT: segments-only ? predictions now in DB; eval/regression operate here
without stitching]
64      -> finally: emit_indexer_report(...) @backend/reporting/__init__.py (? ALWAYS runs
? even on validation failure / exception; NEVER raises; grafts response["reports"])
65      -> select segment ids -> [(start,end)] (+ stitching padding)
66      -> VideoCompiler.compile_highlights(filepath, time_pairs, out_name)
@backend/pipeline/stitcher/compiler.py

```

```

67         per-clip ffmpeg re-encode (libx264/superfast/crf22) -> concat -c copy ->
output/<name>.mp4
68     ...
69     Hybrid segmenter internal 4-phase (yolo_hybrid / tracknet_hybrid / wasb_hybrid ? same
shape):
70     ...
71     P1 chunk video (yolo only; tracknet/wasb = whole-video single pass, chunk_index=0)
72     -> P2 candidate "action windows" from motion/trajectory [STOP: candidates persisted
via db.add_candidate_segment]
73     -> P3 pad clip (ai_handoff_padding) -> ThreadPoolExecutor ->
provider.analyze_video(clip, prompt) ? provider_registry
74     -> P4 db.add_play_segment + db.link_candidate_to_segment
75     ...
76     Pure-AI path `gemini`: no P2; chunk-or-single -> provider -> heavy validate
(clamp/dedup) -> add_play_segment.
77     Legacy `motion`: pixel-diff -> segments + **fabricated** metadata/events (? not ground
truth).
78
79     ### 1B. WASB shuttle substrate (P2 of wasb_hybrid)
80     ...
81     WasbHybridSegmenter.process_video
@backend/pipeline/segmenters/wasb_hybrid.py
82     -> WasbRunner.healthcheck() (gate; not ok -> return [])
@backend/pipeline/detectors/wasb_runner.py
83     -> WasbRunner.run_predict(video, out_dir)
84     stage video + wasb_infer.py into WSL fs -> `wsl bash -c 'source conda.sh && python
wasb_infer.py ...'`
85     -> @backend/pipeline/detectors/wasb_infer.py::run (WSL): sliding windows -> GPU
detector pass (CHECKPOINTED det_raw.jsonl) -> CPU tracker (always full re-run) -> trajectory CSV
86     -> TrackNetRunner.parse_trajectory_csv(csv)
@backend/pipeline/detectors/tracknet_runner.py (shared, re-exported by WasbRunner)
87     -> TrackNetRunner.trajectory_to_action_windows(points, fps, frame_width, chunk_start,
velocity_thresh, merge_gap, min_window_duration, chunk_index) [the model-agnostic WINDOWING
BRAIN]
88     ...
89     (`tracknet_hybrid` identical, swapping WasbRunner->TrackNetRunner; `_probe_fps_width`
fallback (30.0, 1280.0).)
90
91     ### 1C. Calibration / eval flow (NO pipeline run unless told)
92     ...
93     GT CSV -> annotations.load_annotations / calibrate_wasb.load_golden_gts ->
List[Interval]
94     DB predictions -> harness.get_predictions(db, video_id, stage) -> List[Interval] stage
? {candidates, final}
95     -> metrics.match_intervals (greedy temporal-IoU) -> metrics.score -> Score
96     -> harness.evaluate -> EvalReport -> report_to_dict
97     -> batch.aggregate (corpus micro/macro) @backend/eval/batch.py
98     -> regression.regress(_with_provider) vs baseline @backend/eval/regression.py
[CI gate]
99     Calibration: calibrate_wasb.run -> calibration.{select_best, improvement_report,
cross_validate} @backend/eval/calibration.py
100     Distill Gemini->local: calibrate_local (thresholds) OR distill_local (learned
classifier) @backend/eval/{calibrate_local,distill_local}.py
101     Learned segmenter (offline): training.build_training_set(X,y) ->
classifier.LogisticRegression.fit/save
102     Gemini refinement driver (tuning, standalone):
gemini_refine.{refine_window,full_video_rallies} @backend/eval/gemini_refine.py
103     Audio probe (diagnostic only, NOT in live pipeline): audio/e1_probe.run
@backend/eval/audio/e1_probe.py
104     ...
105     Entrypoints: `@run_eval.py` (single video score), `@scripts/run_phase3_eval.py`
(manifest batch, can segment), `@run_regression.py` (baseline gate, exit 1=regression),
`@scripts/nightly_regression.py` (nightly CACHED-trajectory CPU rally-quality tripwire, exit
1=regression/4=stale), `@scripts/nightly_reelcount.py` (nightly FULL-pipeline reel-count +
quality + cost on a GCP GPU VM, emailed), `@backend/eval/calibrate_wasb.py` +

```

```

`@backend/eval/export_wasb_segments.py` (WASB tuning drivers).
106
107     ---
108
109     ## 2. SUBSYSTEMS
110
111     ### 2.0 Orchestration & config
112     - `@backend/main.py` ? Thin CLI + FastAPI bootstrap (orchestrator; modes via FLAGS).
    Static UI mounts `/static` from `backend/api/static`. Logging is initialized with `basicConfig`
    (INFO, timestamped format); backend modules use `logger = logging.getLogger(__name__)` (replaces
    stray prints in hot paths like motion, stitcher, wasb_infer, config loader). User-facing CLI
    summaries may still use `print` for direct output.
113     - `::app` FastAPI; `::db_instance`, `::global_config` (`Optional[AppConfig]`) module
    globals set ONLY in `main()`. ? importing `app` for ASGI without `main()` -> both `None` ->
    every endpoint NPEs.
114     - **ASYNC (#16, PR-1):** `::process_video` `@POST /api/process` -> fast-fail 4xx
    (path/exists/segmenter) -> `db.create_job` -> `_get_executor().submit(_drain_due_jobs)` -> **202
    + job_id**. `::drain_due_jobs` / `::run_claimed_job` (worker loop EXTRACTED to
    `@backend/api/job_worker.py` (#234), re-exported on `backend.main`): claim each runnable job via
    `db.claim_due_job` -> `_execute_processing` -> mark done/failed (a `JobWaiting` parks the job
    via `set_job_waiting` for resume; EXECUTION_POLICY P3 self-scheduling). `::execute_processing`
    = the former sync body (hash -> validate -> add_video -> segment -> `finally:` always
    `emit_indexer_report`, grafts `response["reports"]`; never raises from the report); a RAISING
    segmenter now prune_after-restores pre-run rows before re-raising. `::get_job` `@GET
    /api/jobs/{job_id}` -> `{status, progress, video_id, error, result, reports, ...}` ? job
    `status` = EXECUTION state (`queued|running|done|failed`); the pipeline verdict lives in
    `result["status"]`. `::get_executor` lazy module-global `ThreadPoolExecutor(max_workers=1)`
    (single box + Gemini serial-upload learning; tests monkeypatch it inline). `main()` runs
    `db.fail_stale_jobs()` at startup (orphaned queued/running -> failed). `::compile_highlights`
    `@POST /api/compile`. `::query_play_segments` `@POST /api/query`. `::get_config` `@GET
    /api/config`. `::run_cli_diagnose_clip` `--debug-clip`: lazy `import cv2`, samples  $\pm 3s$  vs
    `indexing.motion_threshold` (dual model/dict read). ? CLI `--video-path` mode still runs the
    pipeline synchronously inline (separate code path, unchanged by #16).
115     - **UPLOAD/DOWNLOAD (PR-2, B2+B3):** the chunked-upload endpoints now live in
    `@backend/api/uploads.py` (#228 ? an `APIRouter` registered via `app.include_router`;
    `uploads.py::upload_cfg` resolves the live `backend.main.global_config` lazily so it reads the
    startup config, and a lazy `import backend.main` avoids an import cycle) ?
    `uploads.py::upload_init` `@POST /api/upload/init` (`{filename, total_size_bytes?}` -> ext
    allowlist + size gate -> `{upload_id}`), `::upload_part` `@PUT /api/upload/{id}/part/{n}` (raw-
    bytes body, STRICTLY sequential n, streams to `partial-<id>` under `security.upload_dir`; per-
    part/total cap breach -> 413 + upload ABORTED), `::upload_complete` `@POST .../complete`
    (`{total_parts}` -> assemble -> never-overwrite rename -> `{path, video_id(MD5), size_bytes}` ?
    client then POSTS /api/process with that path), `::upload_abort` `@DELETE /api/upload/{id}`. ?
    upload state (`uploads.py::uploads` + `::uploads_lock`) is IN-PROCESS ? restart aborts
    partials (client re-inits); per-user partitioning lands with PR-3. `main.py::download_highlight`
    `@GET /api/highlights/{video_id}/{filename}` (download was NOT moved ? still in main.py) ? safe-
    name + `resolve_under_root(output_dir)` then FileResponse stream (the old /api/compile alert
    path was browser-unreachable). Config: `SecurityConfig.{upload_dir='output/uploads',
    max_upload_mb=4096, max_part_mb=95, allowed_upload_extensions=[mp4,mov]}` (95MB: free-tier edge
    proxies cap bodies ~100MB).
116     - **STUDIO / STORAGE endpoints (P1/P2/P3/P4/P6/P8, 2026-07-01):** `::capabilities`
    `@GET /api/capabilities` ? honest box probe; now includes a **`storage` block**
    (`storage_capabilities` in `@backend/storage/capabilities.py`: local/s3/gcs/gdrive
    `{id,kind,label,usable,free_gb}`, secret-free) alongside segmenters/ffmpeg/gpu/deployment.
    `::list_videos` `@GET /api/videos` (the "my videos" list, P6). `::get_source` `@GET
    /api/source/{id}` ? Range-capable FileResponse serving a **frame-identical annotation proxy**
    (D4), source fallback, background-warmed. `::post_annotations` `@POST /api/annotations` ? writes
    the verbatim annotator CSV to `output/labels/<video_id>.rallies.csv` (P8). `::create_asset`
    `@POST /api/assets` (**P4**, `{uploaded_path?|source_uri?, medium_uri, sport?}`) ? stores a
    video into a user-chosen medium (MD5-keyed `storage.put` for an upload?remote medium; register-
    in-place for local; at-rest URI = no copy), registers it (video_id=MD5) so it shows in
    `/api/videos`; **self-host posture only** (403 when `security.allowed_video_root` set); edge-
    auth mirrors `/api/process`. Cost ledger: `::get_job_cost` `@GET /api/jobs/{id}/cost`,
    `::get_video_cost` `@GET /api/videos/{id}/cost`. **Free-space guard (P2):**
    `@backend/utils/freespace.py::require_free_bytes` -> HTTP **507** in `uploads.py::upload_init` +

```



```

the `api/process` remote fetch-to-scratch; the WSL frame-extraction guard now **BLOCKS**
(`detectors/base.py::_wsl_frame_stage_shortfall`, ~15x source).
117     - `::load_config` ? ? legacy thin wrapper returning a **dict**
(`load_app_config(path).model_dump(...)`); kept for transition ? new code imports
`load_app_config`/`AppConfig` directly.
`::ProcessRequest`/`::QueryRequest`/`::CompileRequest`/`::AnnotationsRequest`/`::AssetRequest`
pydantic bodies.
118     - `main()` sets `global_config.output.output_dir` (typed field on
`OutputConfig`); `db_path` = os.path.join(global_config.output.output_dir, output.db_name).
`compile_highlights` reads the SAME `global_config.output.output_dir`, so `--output-dir` is
honored end-to-end (the old write-only `_runtime_overrides` hack is gone).
119     - ? `--segmenter` argparse `choices` frozen at build time from
`segmenter_registry.list_available()` ? a lazily-registered segmenter wouldn't appear. CLI
default_segmenter fallback `motion`.
120     - `@scripts/run_process_and_stitch.py` ? one-shot demo; ? hardcoded paths, interactive
`input()` (blocks automation), skips validation. Config via `backend.config.load_config`
(typed); segmenter, padding, and db path all read from the model (no literal fallbacks).

```

12. user (2026-07-02T04:46:23.637Z)

```

1     """Pluggable storage layer: local (default) + S3/GCS first-class buckets + Google Drive.
2
3     URI-driven (see ``ref.parse_uri``). The package import is cheap ? cloud SDKs are
imported
4     lazily, only when a non-local bucket URI is actually used (install the ``storage-s3`` /
5     ``storage-gcs`` extras). The default ``local`` backend needs no third-party deps and
preserves
6     today's behaviour byte-for-byte; ``gdrive://`` likewise needs none ? it reads a locally
MOUNTED
7     Google Drive (Google Drive for Desktop) as plain files.
8
9     Worker/dispatcher entry points speak URIs via the thin helpers ``fetch`` / ``put`` /
10    ``list_uris`` / ``get_storage``; ``local://`` (and bare paths) make ``fetch`` an
identity
11    passthrough. See ``docs/DECOUPLED_COMPUTE``.
12    """
13
14    from __future__ import annotations
15
16    import os
17    import tempfile
18    from typing import Any, Dict, Iterator, Optional
19
20    from backend.storage.base import StorageBackend
21    from backend.storage.ref import StorageRef, StorageStat, parse_uri, resolve_input_uri
22
23    __all__ = [
24        "StorageBackend",
25        "StorageRef",
26        "StorageStat",
27        "parse_uri",
28        "resolve_input_uri",
29        "get_storage",
30        "fetch",
31        "put",
32        "list_uris",
33        "acquire_local_input",
34        "input_is_local",
35        "resolve_input_ref",
36    ]
37
38
39    def _storage_dict(storage_config: Any) -> Dict[str, Any]:
40        """Coerce a StorageConfig (Pydantic) / dict / None into a plain settings dict."""
41        if storage_config is None:

```

```

42         return {}
43     if hasattr(storage_config, "model_dump"):
44         return storage_config.model_dump() # type: ignore[no-any-return]
45     return dict(storage_config)
46
47
48     def resolve_input_ref(raw: str, storage_config: Any = None) -> StorageRef:
49         """Resolve a user-supplied input (a path or storage URI) into a ``StorageRef`` under
the input
50         config (``input_backend``/``input_root``). Raises ``ValueError`` for an unsupported
URI scheme ?
51         callers at the API/CLI boundary catch that and return a clean client error (never a
500).
52
53         Use ``ref.backend`` to branch local-vs-remote and ``ref.key`` for the resolved LOCAL
path (e.g.
54         ``local://?`` is stripped to the bare path), rather than stat-ing the raw URI
string."""
55         sc = _storage_dict(storage_config)
56         uri = resolve_input_uri(
57             raw,
58             input_backend=sc.get("input_backend", "local"),
59             input_root=sc.get("input_root", ""),
60         )
61         return parse_uri(uri)
62
63
64     def input_is_local(raw: str, storage_config: Any = None) -> bool:
65         """True if ``raw`` resolves (under the input config) to a LOCAL filesystem path ?
i.e. a
66         plain on-disk file the caller can stat/validate directly. False for a bucket / Drive
URI
67         that must be fetched first. Raises ``ValueError`` for an unsupported scheme (same as
68         ``resolve_input_ref``)."""
69         return resolve_input_ref(raw, storage_config).backend == "local"
70
71
72     def acquire_local_input(
73         raw: str, storage_config: Any = None, *, scratch_parent: "str | None" = None
74     ) -> str:
75         """Resolve ``raw`` (a path or storage URI) to a LOCAL file path ready for
processing.
76
77         Local / bare paths are returned UNCHANGED (identity passthrough, no copy ? today's
behaviour
78         byte-for-byte). A bucket / Drive source (``gs://``/``s3://``/``gdrive://``, or a
bare name under a
79         non-local ``input_root``/``input_backend``) is downloaded into a fresh scratch dir
and the local
80         path returned. The scratch copy is left under the system temp dir (a pure
intermediate,
81         regenerable from the source) and relies on OS temp reclamation ? the same pull
pattern the
82         worker's ``cmd_infer`` already ships."""
83         sc = _storage_dict(storage_config)
84         ref = resolve_input_ref(raw, sc)
85         if ref.backend == "local":
86             return (
87                 ref.key
88             ) # identity ? no temp, no copy, byte-identical to passing the path directly
89         scratch = tempfile.mkdtemp(prefix="rally-input-", dir=scratch_parent)
90         return fetch(ref.uri, os.path.join(scratch, ref.name or "in.mp4"), config=sc)
91
92
93     def _backend_class(name: str):

```

```

94     """Lazy-resolve a backend class so cloud SDKs aren't imported until needed."""
95     if name == "local":
96         from backend.storage.local import LocalStorage
97
98         return LocalStorage
99     if name == "s3":
100         from backend.storage.s3 import S3Storage
101
102         return S3Storage
103     if name == "gcs":
104         from backend.storage.gcs import GCSStorage
105
106         return GCSStorage
107     if name == "gdrive":
108         from backend.storage.gdrive import GDriveStorage
109
110         return GDriveStorage
111     if name == "gdrive-api":
112         from backend.storage.gdrive_api import GDriveApiStorage
113
114         return GDriveApiStorage
115     raise ValueError(
116         f"unknown storage backend: {name!r} (expected local|s3|gcs|gdrive|gdrive-api)"
117     )
118
119
120 def get_storage(
121     backend: Optional[str] = None,
122     *,
123     config: Optional[Dict[str, Any]] = None,
124     **overrides,
125 ) -> StorageBackend:
126     """Construct a backend. ``backend`` (or ``config['backend']``, default 'local')
127     selects the implementation; its per-backend settings come from ``config[backend]`` merged
128     with ``overrides``. Settings carry NO secrets ? only endpoints/regions/paths/source-
129     selectors."""
130     cfg = config or {}
131     name = backend or cfg.get("backend", "local")
132     # Per-backend settings live under the backend name; tolerate the hyphen?underscore
133     # hyphenated backend ("gdrive-api") can read a pydantic field that can't contain a
134     # ("gdrive_api"). Hyphen-free names are unaffected (the first lookup wins).
135     settings: Dict[str, Any] = dict(
136         cfg.get(name) or cfg.get(name.replace("-", "_")) or {}
137     )
138     settings.update(overrides)
139     return _backend_class(name)(**settings)
140
141 def fetch(
142     uri: str,
143     dst: "str | os.PathLike[str] | None" = None,
144     *,
145     config: Optional[Dict[str, Any]] = None,
146     overwrite: bool = False,
147 ) -> str:
148     """Resolve ``uri`` to a local file path. For ``local`` this is an identity
149     passthrough (no copy); for s3/gcs it downloads to ``dst``."""
150     ref = parse_uri(uri)
151     return get_storage(ref.backend, config=config).fetch_to_local(
152         ref, dst, overwrite=overwrite
153     )

```

```

153         )
154
155
156     def put(
157         local_path: "str | os.PathLike[str]",
158         uri: str,
159         *,
160         config: Optional[Dict[str, Any]] = None,
161         content_type: Optional[str] = None,
162     ) -> str:
163         """Upload/copy ``local_path`` to ``uri``; return ``uri``."""
164         ref = parse_uri(uri)
165         get_storage(ref.backend, config=config).put(
166             local_path, ref, content_type=content_type
167         )
168         return uri
169
170
171     def list_uris(
172         prefix_uri: str, *, config: Optional[Dict[str, Any]] = None
173     ) -> Iterator[str]:
174         """Yield the URI of every object under ``prefix_uri``."""
175         ref = parse_uri(prefix_uri)
176         for r in get_storage(ref.backend, config=config).list(ref):
177             yield r.uri
178

```

13. user (2026-07-02T04:46:31.298Z)

```

backend\config\models.py:378: # scratch dir before processing (see
``backend.storage.acquire_local_input``).
backend\main.py:646:         local = storage.acquire_local_input(rec["filepath"], getattr(cfg,
"storage", None))
backend\main.py:791:         local_working = storage.acquire_local_input(
backend\main.py:1168:         local_video = storage.acquire_local_input(
backend\main.py:1683:         local_src = storage.acquire_local_input(
backend\main.py:2056:         local_video = storage.acquire_local_input(args.video_path,
storage_cfg)
backend\storage\__init__.py:33:     "acquire_local_input",
backend\storage\__init__.py:72: def acquire_local_input(
docs\PACKAGING_PLAN.md:154:| **CUJ-2** In-UI `gs://` bucket ingest | DONE |
`acquire_local_input` + gcs backend already accept `gs://`; preserved when the user opts into a
cloud storage backend (+ ADC). |
docs\COMPUTE_DECOUPLED_SERVING\README.md:161:- 2026-06-21 ? **M2 landed**
([#280](https://github.com/Khelsutra/badminton-highlight-indexer/pull/280)):
`input_backend`/`input_root` + URI-aware main API/CLI input via
`resolve_input_uri`/`acquire_local_input` (local = identity passthrough, byte-for-byte
unchanged; bucket/Drive fetched to scratch). Unsupported scheme ? clean 400; `local://` accepted
on the resolved key; hosted-mode rejects non-local URIs; source URI kept for DB/report
provenance. Suite 1105 green, cov 84.84%.
tests\test_api_cloud_dispatch.py:114:         m.storage, "acquire_local_input", lambda path,
scfg: str(fake_local)
tests\test_api_cloud_dispatch.py:381:         m.storage, "acquire_local_input", lambda path,
scfg: str(fake_local)
tests\test_api_cloud_dispatch.py:556:         m.storage, "acquire_local_input", lambda path,
scfg: str(fake_local)
tests\test_input_source.py:2: acquire_local_input.
tests\test_input_source.py:19: from backend.storage import acquire_local_input, input_is_local,
resolve_input_ref
tests\test_input_source.py:124: #
----- acquire_local_input
tests\test_input_source.py:125: def test_acquire_local_input_local_identity(tmp_path):
tests\test_input_source.py:129:     assert acquire_local_input(str(p)) == str(p)
tests\test_input_source.py:132: def test_acquire_local_input_relative_local_identity():
tests\test_input_source.py:133:     assert acquire_local_input("rel/clip.mp4") == "rel/clip.mp4"

```

```

tests\test_input_source.py:149:def
test_acquire_local_input_remote_fetches_to_scratch(monkeypatch):
tests\test_input_source.py:152:    out = acquire_local_input(
tests\test_input_source.py:162:def
test_acquire_local_input_input_root_joins_then_fetches(monkeypatch):
tests\test_input_source.py:165:    acquire_local_input("x.mp4", {"input_root":
"gs://bucket/videos"})
tests\test_input_source.py:169:def
test_acquire_local_input_never_fetches_for_local(monkeypatch):
tests\test_input_source.py:174:    assert acquire_local_input("/data/x.mp4") == "/data/x.mp4"
tests\test_input_source.py:176:    acquire_local_input("clip.mp4", {"input_backend":
"local", "input_root": ""})

```

14. user (2026-07-02T04:46:32.234Z)

```

190         outputs_uri = f"{out_root}/outputs/{video_id}/" if video_id else ""
191         report = self._read_json(f"{outputs_uri}report.json") if outputs_uri else {}
192         LOG.info(
193             "cloud-run: job=%s done (video_id=%s rc=%d)", self._job_name, video_id, rc
194         )
195         return ComputeResult(
196             returncode=rc,
197             outputs_uri=outputs_uri,
198             video_id=video_id,
199             report=report,
200             execution=str(execution or ""),
201         )
202
203     # --- storage-backed completion signal (bucket-poll) ---
204     def _read_marker(self, done_uri: str) -> Optional[dict]:
205         """The done-marker dict if the worker has written it, else None ? poll_until
keeps waiting."""
206         ref = parse_uri(done_uri)
207         if not get_storage(ref.backend, config=self._storage_config).exists(ref):
208             return None
209         return (
210             self._read_json(done_uri) or {}
211         ) # empty dict still terminates the poll (present == done)
212
213     def _read_json(self, uri: str) -> dict:
214         # Download to a real local dst FIRST: the cloud backends (gcs/s3) REQUIRE a dst
? a None dst
215         # raises ValueError ? unlike local's identity passthrough. Without an explicit
dst this read
216         # would silently return {} for every gs://|s3:// object (the only backends this
shim targets).
217         try:
218             ref = parse_uri(uri)
219             dst = os.path.join(
220                 tempfile.mkdtemp(prefix="cr-read-"), ref.name or "obj.json"
221             )
222             local = fetch(uri, dst, config=self._storage_config)
223             with open(local, "r", encoding="utf-8") as f:
224                 data: Any = json.load(f)
225             return data if isinstance(data, dict) else {}
226         except Exception:
227             LOG.warning("cloud-run: could not read %s", uri, exc_info=True)
228             return {}
229

```

15. user (2026-07-02T04:46:39.074Z)

```

620         _proxy_inflight.discard(video_id)
621
622

```

```

623     @app.get("/api/source/{video_id}")
624     def get_source(video_id: str, request: Request):
625         """Stream a video for in-UI annotation (P8), Range-capable so the browser can scrub.
Serves a
626         frame-preserving PROXY (STUDIO_MEDIA_LIFECYCLE.md D4 ? downscaled + seekable, labels
still keyed
627         to the master's MD5) once one exists; otherwise serves the source and WARMS a proxy
in the
628         background for next time ? non-blocking, so the first load is never delayed. Proxy
warming is
629         LOCAL-source only (remote-source proxying is a follow-up)."""
630         rec = request.app.state.db.get_video(video_id)
631         if not rec or not rec.get("filepath"):
632             raise HTTPException(status_code=404, detail="Video not found.")
633         cfg = request.app.state.config
634
635         # 1) A ready proxy wins. Bare-name + containment guarded like /api/highlights.
636         try:
637             proxy_name = safe_output_filename(f"{video_id}.mp4")
638             proxy_path = resolve_under_root(os.path.join(_proxies_dir(cfg), proxy_name),
_proxies_dir(cfg))
639         except ValueError:
640             proxy_path = ""
641         if proxy_path and os.path.isfile(proxy_path):
642             return FileResponse(proxy_path, media_type="video/mp4")
643
644         # 2) Otherwise resolve + serve the source.
645         try:
646             local = storage.acquire_local_input(rec["filepath"], getattr(cfg, "storage",
None))
647             is_local = storage.input_is_local(rec["filepath"], getattr(cfg, "storage",
None))
648         except ValueError as e: # unsupported scheme etc. ? a clean client error, never a
500
649             raise HTTPException(status_code=400, detail=str(e)) from e
650         except Exception as e: # noqa: BLE001 - a fetch miss is a bad-gateway, surfaced
(not swallowed)
651             raise HTTPException(status_code=502, detail=f"Could not resolve the source
video: {e}") from e
652         if not os.path.isfile(local):
653             raise HTTPException(status_code=404, detail="Source video file not found.")
654
655         # 3) Warm a proxy in the background for next time (local source only).
656         if is_local:
657             threading.Thread(target=_warm_annotation_proxy, args=(video_id, local, cfg),
daemon=True).start()
658
659         media = "video/quicktime" if local.lower().endswith(".mov") else "video/mp4"
660         return FileResponse(local, media_type=media)
661
662
663     @app.post("/api/annotations")
664     def post_annotations(req: AnnotationsRequest, request: Request):
665         """Persist finished rally labels as the annotator CSV (P8), written verbatim to
666         ``<output_dir>/labels/<video_id>.rallies.csv`` and keyed by the file-MD5
``video_id`` ? so it joins
667         the video by content hash exactly as the training pipeline expects. Byte-compatible
with the VLC
668         plugin + browser extension. A later step (owner-gated) promotes labels into the
durable corpus.
669
670         Edge-auth mirrors /api/process (default-OFF ? no-op locally). ``video_id`` is
validated as a bare
671         filename so an untrusted value can't escape the labels dir (arbitrary write)."""
672         try:

```

```

673         edge_auth.resolve_tenant(request.headers, request.app.state.config)
674     except edge_auth.EdgeAuthError as e:
675         raise HTTPException(status_code=401, detail=f"edge auth: {e}") from e
676
677     cfg = request.app.state.config
678     labels_dir = os.path.join(output_base(cfg), "labels")
679     try:
680         name = safe_output_filename(f"{req.video_id}.rallies.csv")
681     except ValueError as e:
682         raise HTTPException(status_code=400, detail=f"invalid video_id: {e}") from e
683     os.makedirs(labels_dir, exist_ok=True)
684     path = os.path.join(labels_dir, name)
685     try:
686         path = resolve_under_root(path, labels_dir)
687     except ValueError as e:
688         raise HTTPException(status_code=400, detail=str(e)) from e
689     with open(path, "w", encoding="utf-8", newline="") as f:

```

16. user (2026-07-02T04:46:39.920Z)

```

770         validate_input_path(req.uploaded_path, allowed_root=None)
771         src_ref = storage.resolve_input_ref(req.uploaded_path, storage_settings)
772     except ValueError as e:
773         raise HTTPException(status_code=400, detail=str(e)) from e
774     if src_ref.backend != "local" or not os.path.isfile(src_ref.key):
775         raise HTTPException(
776             status_code=404,
777             detail=f"uploaded file not found at '{req.uploaded_path}'",
778         )
779     local_working = src_ref.key
780 else:
781     try:
782         validate_input_path(req.source_uri, allowed_root=None)
783         src_ref = storage.resolve_input_ref(req.source_uri, storage_settings)
784     except ValueError as e:
785         raise HTTPException(status_code=400, detail=str(e)) from e
786     # A remote at-rest source is fetched to scratch to hash it (D3) ? reserve that
scratch (P2).
787     if src_ref.backend != "local":
788         need = _remote_input_size_bytes(src_ref, getattr(cfg, "storage", None))
789         require_free_bytes(tempfile.gettempdir(), need, stage="fetch")
790     try:
791         local_working = storage.acquire_local_input(
792             req.source_uri, getattr(cfg, "storage", None)
793         )
794     except ValueError as e:
795         raise HTTPException(status_code=400, detail=str(e)) from e
796     except Exception as e: # noqa: BLE001 ? a fetch miss is a 502, surfaced not
swallowed
797         raise HTTPException(
798             status_code=502, detail=f"could not resolve source_uri: {e}"
799         ) from e
800     if not os.path.isfile(local_working):
801         raise HTTPException(
802             status_code=404, detail="source_uri did not resolve to a readable file."
803         )
804
805     video_id = compute_video_id(local_working)
806     ext = os.path.splitext(local_working)[1] or ".mp4"
807
808     # 2) Decide the durable location + copy only when needed.
809     copied = False
810     if req.source_uri:
811         stored_uri = req.source_uri # already at rest in its medium ? no copy (D14b)
812         medium_backend = src_ref.backend

```

```

813     else:
814         try:
815             medium_ref = storage.parse_uri(req.medium_uri)
816         except ValueError as e:
817             raise HTTPException(
818                 status_code=400, detail=f"invalid medium_uri: {e}"
819             ) from e
820         medium_backend = medium_ref.backend
821         if medium_ref.backend == "local":
822             # "This computer" ? the upload already lives here; it IS the durable copy
823             (no duplicate).
824             stored_uri = os.path.abspath(local_working)
825         else:
826             dest = _asset_dest_uri(req.medium_uri, video_id, ext)
827             # Free-space: BLOCK a local/Drive-mount destination (D11); a cloud bucket is
828             warn-only.
829             local_dir = _medium_local_dir(medium_ref, storage_settings)
830             if local_dir is not None:
831                 require_free_bytes(
832                     local_dir, os.path.getsize(local_working), stage="store"
833                 )
834             else:
835                 logger.info(
836                     "[ASSETS] cloud medium %s ? capacity not blocked (cost/quota warn
837                     only).",
838                     medium_ref.backend,
839                 )
840             try:
841                 stored_uri = storage.put(local_working, dest, config=storage_settings)
842                 copied = True
843             except ValueError as e:
844                 raise HTTPException(status_code=400, detail=str(e)) from e
845             except Exception as e: # noqa: BLE001 ? a put failure is surfaced, never
846                 silently dropped
847                 logger.error(
848                     "[ASSETS] put failed video_id=%s dest=%s: %s",
849                     video_id,
850                     dest,
851                     e,
852                     exc_info=True,
853                 )
854                 raise HTTPException(
855                     status_code=502, detail=f"could not store into the medium: {e}"
856                 ) from e
857
858     # 3) Register the asset (keyed by MD5) so it appears in GET /api/videos ("my
859     videos").
860     if not request.app.state.db.add_video(video_id, stored_uri, "stored", []):
861         raise HTTPException(status_code=500, detail="failed to register the asset.")
862     logger.info(
863         "[ASSETS] stored video_id=%s medium=%s uri=%s copied=%s owner=%s",
864         video_id,
865         medium_backend,
866         stored_uri,
867         copied,
868         owner_id,
869     )
870     return JSONResponse(
871         status_code=201,
872         content={
873             "video_id": video_id,
874             "medium": medium_backend,
875             "stored_uri": stored_uri,
876             "copied": copied,
877             "sport": req.sport,

```



```

873         "status": "stored",
874     },
875 )
876
877
878 def _ingest_remote_segments(
879     db, video_id: str, outputs_uri: str, storage_config: dict, *, max_rows: int = 20000
880 ) -> int:
881     """Read a cloud job's ``intervals.csv`` (start,end,shot_count,ending_reason) from
882     the bucket and
883     persist each rally via the SAME ``db.add_play_segment`` the in-process segmenters
884     call ? so
885     ``/api/query`` + ``/api/compile`` read the rows with zero knowledge of where they
886     were produced.
887     Returns the count ingested. Mirrors ``cloud_run._read_json``'s fetch-to-a-real-dst
888     pattern (cloud
889     backends REQUIRE a dst). Hardened against a hostile/corrupt worker output: a
890     malformed row is
891     skipped, non-finite / overflowing values are rejected, the loop is bounded, and the
892     temp dir is
893     cleaned. May raise (missing/unreadable CSV, fetch error) ? the caller treats that as
894     a job failure
895     and prunes any partial rows (destroy-AFTER-confirm)."""
896     uri = f"{outputs_uri.rstrip('/')}/intervals.csv"

```

17. user (2026-07-02T04:46:46.438Z)

```

1130
1131
1132 def _execute_processing(
1133     config=None, db=None, req=None, progress=None
1134 ) -> Dict[str, Any]:
1135     """The full hash ? validate ? segment ? report pipeline for one video.
1136
1137     Accepts explicit ``config``/``db`` (the new app-factory pattern) OR reads the
1138     module-level ``global_config``/``db_instance`` (backward-compatible with existing
1139     tests that pass only ``req``). Route handlers always pass explicit params; older
1140     test code that calls ``_execute_processing(req)`` still works.
1141
1142     ``progress`` is an optional callable(stage: str) used to surface coarse job
1143     progress.
1144     Raises on unexpected internal errors ? the caller records those as a job failure
1145     """
1146     # Backward-compatible shim: if called as _execute_processing(req) (old pattern),
1147     # shift positional args ? req arrives in position 0 (config), progress in position 1
1148     (db).
1149     if config is not None and not isinstance(config, AppConfig):
1150         # Old signature: _execute_processing(req, progress=None)
1151         if req is None and db is not None and isinstance(db, ProcessRequest):
1152             req = db
1153             db = None
1154         elif req is None and isinstance(config, ProcessRequest):
1155             req = config
1156             config = None
1157     # Resolve config/db from module globals when not passed explicitly
1158     if config is None:
1159         config = global_config
1160     if db is None:
1161         db = db_instance
1162
1163     notify = progress or (lambda stage: None)
1164
1165     notify("hashing")
1166     # M2: resolve the input to a LOCAL path. A bare/local path is returned unchanged
1167     (identity ?

```

```

1165         # byte-identical to today)); a bucket/Drive URI is fetched to scratch first. The
original
1166         # req.video_path (the source URI) is kept for the DB record + report; only the file-
reading
1167         # consumers (hash / validators / segmenter) use the resolved local path.
1168         local_video = storage.acquire_local_input(
1169             req.video_path, getattr(config, "storage", None)
1170         )
1171         video_id = compute_video_id(local_video)
1172         was_reingest = db.get_video(video_id) is not None
1173
1174         # 1. Run pluggable validators (with-context variant surfaces ffprobe_meta for the
report)
1175         notify("validating")
1176         logger.info(f"Running validations for {req.video_path}...")
1177         val_results, val_context = registry.run_all_with_context(local_video, config)
1178         failures = [r for r in val_results if not r.passed]
1179
1180         # typed AppConfig: attribute access for the default segmenter (with safe fallback)
1181         default_seg = getattr(
1182             getattr(config, "indexing", None), "default_segmenter", "motion"
1183         )
1184         seg_name = req.segmenter or default_seg
1185         segmenter_actual = None
1186         segmentation_ran = False
1187         # Compute-path PROVENANCE (the validation method): which backend actually ran the
DETECT ?
1188         # "in_process" once the local segmenter is reached, set on the response as
"cloud_run" by
1189         # _maybe_dispatch_remote when it dispatches. None ? nothing ran (validation/config
fail).
1190         compute_actual: Optional[str] = None
1191         response: Optional[Dict[str, Any]] = None
1192         try:
1193             if failures:
1194                 # add_video is an UPSERT (no cascade) ? a failed re-validation only updates
1195                 # status; it no longer destroys the video's prior good segments.
1196                 db.add_video(
1197                     video_id,
1198                     req.video_path,
1199                     "failed",
1200                     [r.model_dump() for r in val_results],
1201                 )
1202                 response = {
1203                     "video_id": video_id,
1204                     "status": "failed",
1205                     "message": "Video failed validation checks.",
1206                     "results": [r.model_dump() for r in val_results],
1207                 }
1208                 return response
1209
1210         # 2. Run segmenter & indexer
1211         logger.info("[API] Video passed checks. Segmenting and indexing...")
1212         db.add_video(
1213             video_id, req.video_path, "processed", [r.model_dump() for r in val_results]
1214         )
1215
1216         segmenter_cls = segmenter_registry.get(seg_name)
1217         if not segmenter_cls:
1218             # Submit-time validation already 400s unknown names; this is the defensive
1219             # in-worker path (e.g. config default pointing at an unregistered
segmenter).
1220             available = list(segmenter_registry.list_available().keys())
1221             response = {
1222                 "video_id": video_id,

```

```

1223         "status": "failed",
1224         "message": f"Segmenter '{seg_name}' not found. Available segmenters:
{available}",
1225         "results": [r.model_dump() for r in val_results],
1226         "play_segments": [],
1227     }
1228     return response
1229
1230     logger.info(f"[API] Using segmenter: {seg_name}")
1231     notify(f"segmenting ({seg_name})")
1232     segmenter_actual = seg_name
1233     # Destroy-AFTER-confirm: remember the pre-run rows, then keep new + drop old
1234     only
1235     # if the run yields segments; otherwise drop the new partial rows and keep the
1236     old.
1237     play_wm, cand_wm = db.segment_watermarks(video_id)
1238     # COMPUTE_DECOUPLED_SERVING M6 ? API: with deployment.compute_target=cloud_run,
1239     run the heavy
1240     # DETECT on a Cloud Run GPU Job + ingest its rallies HERE so query/compile are
1241     unchanged.
1242     # DEFAULT-OFF: returns None for every non-cloud target ? the in-process
1243     segmenter below runs
1244     # byte-identically. (player_pool/max_frames aren't threaded to the worker
1245     `infer` CLI yet, so
1246     # they're dropped on the cloud path in this slice ? a documented follow-up.)
1247     remote = _maybe_dispatch_remote(
1248         config, db, req, video_id, seg_name, val_results, play_wm, cand_wm, notify
1249     )
1250     if remote is not None:
1251         response, segmentation_ran = (
1252             remote # `ran` = did the cloud job actually execute (report)
1253         )
1254         return response # the cloud branch already stamped response["compute"]
1255     (path=cloud_run)
1256     # In-process from here on ? record the actual path for the provenance stamp in
1257     `finally`.
1258     compute_actual = "in_process"
1259     segmenter = segmenter_cls(db, config)
1260     try:
1261         result = segmenter.process_video(
1262             video_id, local_video, req.player_pool, req.max_frames
1263         )
1264     except Exception:
1265         # A raising segmenter must not leave half-written rows: restore the pre-run
1266         # state (same destroy-after-confirm contract as the Failure branch), then
1267         let
1268         # the job worker record the failure.
1269         db.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)
1270         raise
1271     segmentation_ran = True
1272
1273     if isinstance(result, Failure):
1274         logger.error(f"[API] Segmenter failed: {result.message}")
1275         db.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)
1276         response = {
1277             "video_id": video_id,
1278             "status": "failed",
1279             "message": f"Segmenter '{seg_name}' failed: {result.message}",
1280             "results": [r.model_dump() for r in val_results],
1281             "play_segments": [],
1282         }
1283     }
1284     return response
1285
1286     segments = result.value if hasattr(result, "value") else result
1287     notify("finalizing")

```

```

1278         _finalize_reingest(db, video_id, segments, play_wm, cand_wm)
1279
1280         response = {
1281             "video_id": video_id,
1282             "status": "success",
1283             "message": f"Successfully indexed {len(segments)} play segments using
1284             '{seg_name}' segmenter.",
1285             "results": [r.model_dump() for r in val_results],
1286             "play_segments": segments,
1287         }
1288         return response
1289     finally:
1290         # Compute-path PROVENANCE stamp (the validation method): the cloud branch
1291         already set
1292         # response["compute"] (path=cloud_run + the Cloud Run execution); for every
1293         other path stamp
1294         # the actual path here so the job result PROVES which backend ran (in_process /
1295         not_run) and
1296         # what the picker REQUESTED ? a silent fallback (requested cloud, ran local) is
1297         then visible.
1298         if isinstance(response, dict):
1299             response.setdefault(
1300                 "compute",
1301                 {
1302                     "path": compute_actual or "not_run",
1303                     "requested": getattr(req, "compute", None),
1304                 },
1305             )
1306         prov = response.get("compute", {}) if isinstance(response, dict) else {}
1307         logger.info(
1308             "[COMPUTE] video_id=%s requested=%s actual=%s dispatched=%s status=%s",
1309             video_id,
1310             getattr(req, "compute", None),
1311             prov.get("path"),
1312             prov.get("dispatched"),
1313             (response or {}).get("status") if isinstance(response, dict) else None,
1314         )
1315         # Always emit the per-video observability report (next to the video, or to
1316         # report.output_dir). Never raises. Surface the paths in the response. (The
1317         compute-path
1318         # proof lives on the job result + the [COMPUTE] log, not the report ?
1319         obsreport's record_run
1320         # doesn't surface arbitrary run signals.)
1321         paths = emit_indexer_report(
1322             db=db,
1323             video_id=video_id,
1324             # Resolved LOCAL path, not the raw URI: obsreport writes the report next to
1325             the
1326             # video via os.path.abspath(video_path), and a remote URI (gs://, gcs://,
1327             local://)
1328             # collapses to a literal scheme dir at the repo root (see #397). The source
1329             URI is
1330             # still preserved in DB provenance via db.add_video.
1331             video_path=local_video,
1332             config=config,
1333             val_results=val_results,
1334             val_context=val_context,
1335             segmenter_requested=req.segmenter,
1336             segmenter_actual=segmenter_actual,
1337             was_reingest=was_reingest,
1338             segmentation_ran=segmentation_ran,
1339         )
1340         if paths and isinstance(response, dict):
1341             response["reports"] = paths
1342

```

```

1333
1334     def _remote_input_size_bytes(input_ref, storage_config) -> Optional[int]:
1335         """Best-effort size (bytes) of a remote storage source, or ``None`` if unknowable.
1336
1337         Used only to pre-reserve scratch for a fetch-to-local (P2). Never raises: a stat
that needs a
1338         missing cloud SDK, the network, or a credential returns ``None`` so the free-space
guard degrades
1339         to a no-op rather than blocking a job it cannot measure."""
1340         try:
1341             sc = (
1342                 storage_config.model_dump()
1343                 if hasattr(storage_config, "model_dump")
1344                 else (storage_config or {})
1345             )
1346             backend = storage.get_storage(input_ref.backend, config=sc)
1347             return int(backend.stat(input_ref).size)
1348         except Exception as e: # noqa: BLE001 - measuring is best-effort; inability must
not block
1349             logger.debug(
1350                 "free-space: could not stat remote source %s: %s", input_ref.uri, e
1351             )
1352             return None
1353
1354
1355     @app.post("/api/process", status_code=202)
1356     def process_video(req: ProcessRequest, request: Request):
1357         """Submit a processing job (202 + job_id). Poll GET /api/jobs/{job_id} for state.
1358
1359         Fast-fail checks (bad path, missing file, unknown segmenter) happen here so the

```

18. user (2026-07-02T04:46:53.566Z)

```

1640         content={"job_id": job_id, "status": "queued", "poll": f"/api/jobs/{job_id}"},
1641     )
1642
1643
1644     def _execute_compile(cfg, db, job: Dict[str, Any], progress=None) -> Dict[str, Any]:
1645         """Worker handler for a ``kind='compile'`` job (#329): re-resolve the segments,
stitch the reel,
1646         and return the {status, filepath, download_url} result the SPA polls for. Returns a
1647         ``status='failed'`` result for an EXPECTED problem (missing video / no segments / a
stitch error)
1648         so the SPA always gets a clean verdict instead of a worker-crash error."""
1649         notify = progress or (lambda stage: None)
1650         params = json.loads(job.get("params") or "{}")
1651         video_id = job.get("video_id") or ""
1652         segment_ids = params.get("segment_ids") or []
1653         locale = params.get("locale")
1654
1655         notify("resolving")
1656         try:
1657             video, kept, out_name = _resolve_compile(
1658                 db, cfg, video_id, segment_ids, params.get("output_filename", "")
1659             )
1660         except _CompileError as e:
1661             return {"status": "failed", "message": e.detail, "video_id": video_id}
1662
1663         stitching = getattr(cfg, "stitching", None) or StitchingConfig()
1664         output_dir = (
1665             getattr(getattr(cfg, "output", None), "output_dir", "output") or "output"
1666         )
1667         # Localize the branding watermark to the UI locale the SPA rendered in (recorded on
the job);
1668         # unknown/absent locale keeps the configured default unchanged.

```

```

1669     localized_watermark = localize_watermark(stitching.watermark, locale)
1670     compiler = VideoCompiler(
1671         output_dir,
1672         begin_padding=stitching.begin_padding,
1673         end_padding=stitching.end_padding,
1674         watermark=localized_watermark,
1675     )
1676     time_pairs = [(s["start_time"], s["end_time"]) for s in kept]
1677
1678     notify("stitching")
1679     try:
1680         # The stored `filepath` is the SOURCE ref (gs://, s3://, gdrive://, or a local
1681         # path). Resolve to
1682         # a LOCAL file the stitcher can read ? identity for a local path, fetch-to-
1683         # scratch for a URI (the
1684         # same seam the process path uses); without it, compiling a bucket-ingested
1685         # video can't open gs://.
1686         local_src = storage.acquire_local_input(
1687             video["filepath"], getattr(cfg, "storage", None)
1688         )
1689         out_path = compiler.compile_highlights(local_src, time_pairs, out_name)
1690     except Exception as e: # noqa: BLE001 ? a stitch/ffmpeg failure is a clean job
1691         FAILURE, not a crash
1692         logger.error(
1693             "[compile] stitching failed for video_id=%s: %s", video_id, e, exc_info=True
1694         )
1695     return {
1696         "status": "failed",
1697         "message": f"Highlight stitching failed: {e}",
1698         "video_id": video_id,
1699     }
1700
1701     # `filepath` is server-local (not browser-reachable); also return the download URL
1702     # so the SPA needn't
1703     # hand-build it (P6 / PER_RALLY_SELECTION.md §4.3).
1704     download_url = f"/api/highlights/{quote(video_id)}/{quote(out_name)}"
1705     notify("finished")
1706     return {
1707         "status": "success",
1708         "filepath": out_path,
1709         "download_url": download_url,
1710         "video_id": video_id,
1711     }
1712
1713 # --- CLI Debug Utility & Orchestration ---
1714 def run_cli_diagnose_clip(video_path: str, timestamp: float):
1715     """CLI debugging utility to examine segment properties around a timestamp."""
1716     print("=" * 60)
1717     print(f"DEBUGGING VIDEO CLIP: {video_path} at {timestamp}s")
1718     print("=" * 60)
1719
1720     cfg = load_config() # legacy wrapper still works, returns dict for now
1721
1722     # 1. Validation check diagnostics
1723     print("[DIAGNOSTIC] Running validation framework...")
1724     results = registry.run_all(video_path, cfg)
1725     for r in results:
1726         status_symbol = "[PASS]" if r.passed else "[FAIL]"
1727         print(f"  {status_symbol} {r.validator_name}: {r.message}")
1728         if r.details:
1729             print(f"    Details: {r.details}")
1730
1731     # 2. Local segment analysis around the timestamp (sampling +/- 5 seconds)
1732     print("-" * 60)

```

```

1729     print("[DIAGNOSTIC] Analyzing local motion around the timestamp...")
1730     # Lazy import: cv2 is NOT installed in the LIGHT tier (torch-free default)
1731     # and this --debug-clip path is CLI-only, not on the served API path.
1732     import cv2
1733
1734     cap = cv2.VideoCapture(video_path)
1735     if not cap.isOpened():
1736         print("[FAIL] OpenCV could not open video.")
1737         return
1738
1739     fps = cap.get(cv2.CAP_PROP_FPS)
1740     total_frames = cap.get(cv2.CAP_PROP_FRAME_COUNT)
1741
1742     start_frame = max(0, int((timestamp - 3.0) * fps))
1743     end_frame = min(int(total_frames), int((timestamp + 3.0) * fps))
1744
1745     # Measure frame-by-frame changes in sample region
1746     cap.set(cv2.CAP_PROP_POS_FRAMES, start_frame)
1747     prev_gray = None
1748     motions = []
1749
1750     for _ in range(start_frame, end_frame):
1751         ret, frame = cap.read()
1752         if not ret:
1753             break
1754         gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
1755         gray = cv2.GaussianBlur(gray, (21, 21), 0)
1756
1757         if prev_gray is not None:
1758             diff = cv2.absdiff(prev_gray, gray)
1759             _, thresh = cv2.threshold(diff, 25, 255, cv2.THRESH_BINARY)
1760             motion_ratio = cv2.countNonZero(thresh) / (
1761                 thresh.shape[0] * thresh.shape[1]
1762             )
1763             motions.append(motion_ratio)
1764         prev_gray = gray
1765
1766     cap.release()
1767
1768     if motions:
1769         avg_motion = sum(motions) / len(motions)

```

19. user (2026-07-02T04:46:59.420Z)

```

1960         return
1961         run_cli_diagnose_clip(args.video_path, args.debug_clip)
1962         return
1963
1964     # P2b (PACKAGING_PLAN.md gate 08): batteries-included first run for the friendly
1965     `--app`
1966     # launcher. Seed a truly-local, torch-free LIGHT default config (detector 'motion' ?
1967     offline,
1968     # no API key, no GPU, no cloud bills) at the resolved app-home
1969     (default_config_path(), which
1970     # roots under RALLY_APP_HOME when the packaged shortcut pins it) BEFORE
1971     load_app_config reads
1972     # it ? so a fresh `pip`/`uv tool install` + `indexer --app` just works. NEVER
1973     overwrites an
1974     # existing config (a user's edits / the P2c wizard's choices are preserved). Scoped
1975     to --app
1976     # (the packaged app mode); `--server` and the CLI single-shot are byte-identical (no
1977     seeding).
1978     if args.app:
1979         seeded_path, seeded = write_light_default_config()
1980         if seeded:

```

```

1974         print(
1975             f"[APP] First run: seeded a truly-local default config at {seeded_path}"
1976             "
1977             f"(detector 'motion' ? offline, no API key or GPU needed). "
1978             f"Use the in-app setup to switch to Gemini."
1979         )
1980     # Initialize Global Config and DB
1981     global global_config, db_instance
1982     global_config = load_app_config() # returns typed AppConfig
1983     _enforce_startup_or_exit(
1984         global_config
1985     ) # M5: fail fast on an incoherent ACTIVE profile
1986
1987     # The CLI --output-dir overrides the configured output directory. It now lives
1988     # on the typed model (OutputConfig.output_dir), so every consumer ? including
1989     # /api/compile ? reads the same value instead of a write-only runtime hack.
1990     # P0a: a RELATIVE output dir roots under the app-home (RALLY_APP_HOME); with the env
1991     # unset, resolve_output_dir("output") == "output" (CWD-relative) ? byte-identical to
before.
1992     global_config.output.output_dir = resolve_output_dir(args.output_dir)
1993     os.makedirs(global_config.output.output_dir, exist_ok=True)
1994
1995     db_path = os.path.join(
1996         global_config.output.output_dir, global_config.output.db_name
1997     )
1998
1999     # P0e: refuse a SECOND instance against the same DB BEFORE the destructive stale-job
sweep below.
2000     # The job worker is in-process, so a concurrent process opening this DB would sweep
THIS one's
2001     # in-flight jobs to 'failed'. The OS lock (held for the process lifetime; auto-
released on exit
2002     # or crash) is keyed to db_path, so a different --output-dir/DB still runs
concurrently.
2003     global _instance_lock
2004     try:
2005         _instance_lock = acquire_lock(db_path + ".lock")
2006     except OSError as exc:
2007         # The lock path isn't writable (read-only dir / perms). Don't crash and don't
falsely claim a
2008         # second instance ? warn and proceed WITHOUT the guard (a lock-infra failure
must not block a
2009         # legitimate single instance).
2010         logger.warning(
2011             "[JOBS] Could not create the single-instance lock at %s.lock (%s);
proceeding "
2012             "WITHOUT the second-instance guard.",
2013             db_path,
2014             exc,
2015         )
2016         _instance_lock = None
2017     else:
2018         if _instance_lock is None:
2019             print(
2020                 f"[FATAL] Another instance is already using this database ({db_path}).
Refusing to "
2021                 f"start ? a second instance would corrupt in-flight job state. Stop the
other one, or "
2022                 f"use a different --output-dir (RALLY_APP_HOME) to run an independent
instance."
2023             )
2024             sys.exit(1)
2025
2026     db_instance = Database(db_path)

```



```

2027
2028 # Crash recovery (#16): jobs left queued/running by a previous process are dead ?
2029 # the executor is in-process. Mark them failed so pollers see a terminal state.
2030 swept = db_instance.fail_stale_jobs()
2031 if swept:
2032     logger.warning(
2033         f"[JOBS] Swept {swept} stale job(s) from a previous run ? failed."
2034     )
2035
2036 # 01/P11: sweep Gemini File-API uploads orphaned by a hard process death (the per-
call
2037 # _safe_delete covers normal/except paths, but a SIGKILL/OOM mid-upload leaves a
clip on
2038 # Google's servers forever). Best-effort, video-only, grace-gated; never blocks
startup.
2039 _maybe_sweep_gemini_orphans(global_config)
2040
2041 # CLI processing mode (no web UI needed). --app/--server always run the server,
never CLI single-shot.
2042 if args.video_path and not args.server and not args.app:
2043     print(f"[CLI] Processing video file: {args.video_path}")
2044     # M2: bare/local path = identity (today's behaviour); a bucket/Drive URI is
fetched to scratch.
2045     # The local-existence fast-fail only applies to a local input (stat the RESOLVED
key, not the
2046     # raw URI); a remote fetch fails loudly. An unsupported scheme is a clean error,
not a crash.
2047     storage_cfg = getattr(global_config, "storage", None)
2048     try:
2049         input_ref = storage.resolve_input_ref(args.video_path, storage_cfg)
2050     except ValueError as e:
2051         print(f"[FAIL] {e}")
2052         return
2053     if input_ref.backend == "local" and not os.path.exists(input_ref.key):
2054         print(f"[FAIL] Video file not found: {args.video_path}")
2055         return
2056     local_video = storage.acquire_local_input(args.video_path, storage_cfg)
2057
2058     video_id = compute_video_id(local_video)
2059     was_reingest = db_instance.get_video(video_id) is not None
2060
2061     # Validate (with-context variant surfaces ffprobe_meta for the report)
2062     print("[CLI] Validating...")
2063     val_results, val_context = registry.run_all_with_context(
2064         local_video, global_config
2065     )
2066     failures = [r for r in val_results if not r.passed]
2067
2068     # Save video status
2069     status = "failed" if failures else "processed"
2070     db_instance.add_video(
2071         video_id, args.video_path, status, [r.model_dump() for r in val_results]
2072     )
2073
2074     seg_name = args.segmenter or getattr(
2075         getattr(global_config, "indexing", None), "default_segmenter", "motion"
2076     )
2077     segmenter_actual = None
2078     segmentation_ran = False
2079     try:
2080         print("\n" + "=" * 40)
2081         print("VALIDATION SUMMARY")
2082         print("=" * 40)
2083         for r in val_results:
2084             status_symbol = "[PASS]" if r.passed else "[FAIL]"

```

```

2085         print(f"{status_symbol} {r.validator_name}: {r.message}")
2086     print("=" * 40 + "\n")
2087
2088     if failures:
2089         print("[FAIL] Video failed checks. Indexing halted.")
2090         return
2091
2092     # Index
2093     segmenter_cls = segmenter_registry.get(seg_name)
2094     if not segmenter_cls:
2095         print(f"[FAIL] Segmenter '{seg_name}' not found.")
2096         return
2097
2098     print(f"[CLI] Indexing play segments using segmenter: {seg_name}...")
2099     segmenter_actual = seg_name
2100     play_wm, cand_wm = db_instance.segment_watermarks(video_id)
2101     segmenter = segmenter_cls(db_instance, global_config)
2102     result = segmenter.process_video(
2103         video_id, local_video, max_frames=args.max_frames
2104     )
2105     segmentation_ran = True
2106
2107     if isinstance(result, Failure):
2108         print(f"[FAIL] Segmenter '{seg_name}' failed: {result.message}")
2109         print("[FAIL] Indexing halted for video.")
2110         segments = []
2111     else:
2112         segments = result.value if hasattr(result, "value") else result
2113         print(
2114             f"[SUCCESS] Process complete. Indexed {len(segments)} play segments
inside database: {db_path}"
2115         )
2116         # Destroy-AFTER-confirm: keep new on success, restore prior on
empty/failure.
2117         _finalize_reingest(db_instance, video_id, segments, play_wm, cand_wm)
2118     finally:
2119         # Always emit the per-video observability report. Never raises.
2120         paths = emit_indexer_report(
2121             db=db_instance,
2122             video_id=video_id,
2123             # Resolved LOCAL path, not the raw URI (see #397) ? a remote URI would
write a
2124             # literal scheme dir (gcs:/, gs:/, local:/) at the repo root. Source URI
is kept
2125             # in DB provenance via db.add_video.
2126             video_path=local_video,
2127             config=global_config,
2128             val_results=val_results,
2129             val_context=val_context,
2130             segmenter_requested=args.segmenter,
2131             segmenter_actual=segmenter_actual,
2132             was_reingest=was_reingest,
2133             segmentation_ran=segmentation_ran,
2134         )
2135     if paths:
2136         print(f"[REPORT] Technical report : {paths.get('technical_md')}")
2137         print(f"[REPORT] Customer summary : {paths.get('customer_md')}")
2138         print(f"[REPORT] Machine JSON      : {paths.get('json')}")
2139     return
2140
2141     # Start Local Server Mode
2142     if args.app or args.server or not args.video_path:
2143         if args.app:
2144             # P2: the friendly launcher ? always localhost-only (use --server for
network exposure),

```

```

2145         # auto-pick a free port if --port is busy, and open the bundled UI once the
server is up.
2146         bind_host = "127.0.0.1"
2147         port = _resolve_app_port(bind_host, args.port)
2148         if port != args.port:
2149             print(f"[APP] Port {args.port} is busy ? using {port} instead.")
2150         url = f"http://{bind_host}:{port}/"
2151         # The browser-open health poll sends Host: 127.0.0.1 ? it relies on loopback
staying in
2152         # the Host allowlist (the default; don't set RALLY_ALLOWED_HOSTS to a non-
loopback-only set
2153         # with --app, or the poll 400s and the browser just won't auto-open ? the
server still runs).
2154         print(
2155             f"[APP] Starting KhelSutra ? opening {url} in your browser "
2156             f"(set RALLY_NO_BROWSER=1 to skip). Press Ctrl+C to stop."
2157         )
2158         threading.Thread(
2159             target=_open_browser_when_ready,
2160             args=(url, url + "api/health"),
2161             daemon=True,
2162         ).start()
2163     else:
2164         bind_host = _resolve_bind_host(
2165             global_config, args.host
2166         ) # P0c: config/profile-driven (default 127.0.0.1)
2167         port = args.port
2168         print(
2169             f"[SERVER] Starting local FastAPI server on
http://{bind_host}:{port}..."
2170         )
2171         if bind_host not in ("127.0.0.1", "localhost", "::1"):
2172             print(
2173                 f"[SERVER] ? Binding {bind_host} exposes the API on the network ?
ensure an auth "
2174                 f"gate (edge_auth_enabled / a reverse proxy) is in front, and set "
2175                 f"RALLY_ALLOWED_HOSTS to your public host(s) (the Host-header guard
defaults to "
2176                 f"loopback-only and will 400 forwarded requests otherwise)."
2177             )
2178         # Resume self-scheduled work orphaned by the restart (EXECUTION_POLICY.md P3):
the
2179         # in-process wake-timer for any parked 'waiting' job died with the old process,
so
2180         # kick one drain ? it reclaims due jobs and re-arms the wake-timer for the rest.
2181         _get_executor().submit(_drain_due_jobs, global_config, db_instance)
2182         # Stamp shared state on the app so route handlers access it via
request.app.state
2183         # instead of the Optional module-level globals (CODE_CLEANUP_AUDIT P5 / D2).
2184         create_app(global_config, db_instance)
2185         uvicorn.run(app, host=bind_host, port=port)
2186
2187
2188     if __name__ == "__main__":
2189         # `python -m backend.main` runs THIS file as `__main__`, but `job_worker`/`uploads`
resolve the
2190         # shared mutable state (`db_instance`, `global_config`) via `import backend.main` ?
a SEPARATE
2191         # module object. If `main()` runs here it configures `__main__`'s globals, leaving
2192         # `backend.main.db_instance` = None; the async-job drain then crashes silently on
2193         # `None.claim_due_job()` and every submitted job hangs in 'queued' forever. Dispatch
through
2194         # `backend.main` so `main()` configures the module everyone actually reads. (No
recursion:
2195         # imported as `backend.main`, `__name__` != "__main__", so this block doesn't re-

```

```

fire.)
2196         import backend.main as _bm
2197
2198         _bm.main()
2199

```

20. user (2026-07-02T04:47:06.008Z)

```

1      """Async job worker ? drain/wake loop (DEFERRED_HARDENING_PLAN #16, EXECUTION_POLICY.md
P3).
2
3      Extracted verbatim from backend.main to shrink the orchestrator; behavior is unchanged.
4      The worker DRAINS the jobs table via `claim_due_job` and honours WAIT_AND_RESUME: a job
5      that raises `JobWaiting` is parked (status 'waiting' + `next_poll_at` + progress) and
6      re-claimed when due ? with no central scheduler (the DB is the clock).
7
8      The monkeypatch contract these functions rely on (and that tests depend on):
9      every cross-function seam ? `db_instance`, `global_config`, `_execute_processing`,
10     `_arm_wake_timer`, `_get_executor`, `_MAX_DRAIN_WAKE_SEC` ? is resolved THROUGH the
11     `backend.main` module object at call-time (`import backend.main as _main`), never as a
12     bare local. `backend.main` re-exports these names, so `monkeypatch.setattr(backend.main,
13     ...) ` still intercepts the inter-function calls (e.g. a parked job arms the *patched*
14     `_arm_wake_timer`, so no real daemon timer leaks in tests). Importing `backend.main`
lazily
15     inside the functions also avoids the circular import at module load (main.py imports
this
16     module; this module must not import main at import time).
17     """
18
19     import logging
20     import threading
21     import traceback
22     from concurrent.futures import ThreadPoolExecutor
23     from typing import Any, Dict, Optional
24
25     logger = logging.getLogger(__name__)
26
27     # Async job executor (DEFERRED_HARDENING_PLAN #16). max_workers=1 on purpose: a single
28     # self-hosted box serializes GPU work, and the Gemini provider is rate-fragile under
29     # parallel uploads (2026-06-10 ops learning: serial uploads + backoff). The segmenters
30     # already parallelize their AI handoff internally via their own pools.
31     #
32     # ? 02 RISK (CODE_CLEANUP_AUDIT $3c): a single worker means ONE stuck or hung job
33     # (e.g. hung Gemini generate, hung WSL GPU call, hung Cloud Run poll) blocks ALL
34     # subsequent jobs ? they stay 'queued' forever. Mitigations already in place:
35     #   - ExecutionPolicy op_timeout_sec kills hung generate/processing calls
36     #   - detector timeout_sec kills hung WSL/GPU calls
37     #   - Cloud Run poll has a max_wait_sec bound
38     # A future slice should add a per-job wall-clock timeout (the executor itself
39     # cannot enforce timeouts on the Thread).
40     _job_executor: Optional[ThreadPoolExecutor] = None
41
42
43     def _get_executor() -> ThreadPoolExecutor:
44         """Lazy module-global executor; tests monkeypatch this for inline execution."""
45         global _job_executor
46         if _job_executor is None:
47             _job_executor = ThreadPoolExecutor(
48                 max_workers=1, thread_name_prefix="jobworker"
49             )
50         return _job_executor
51
52
53     class JobWaiting(Exception):
54         """Raised inside the pipeline to SELF-SCHEDULE a job (EXECUTION_POLICY.md P3,

```

```

55     WAIT_AND_RESUME). It means "I can't finish now ? park me and resume in `delay_sec`",
56     NOT a failure. The worker persists `next_poll_at` + `progress` (via
`set_job_waiting`),
57     releases the thread, and re-claims the job when due (`claim_due_job`). No work to
date
58     is the segmenter's `_execute_processing`, run UNDER the destroy-after-confirm
contract.
59
60     The wiring is additive: the production segmenter path never raises this today (zero
61     behaviour change), so a job runs exactly as before. The hook is what a later slice
62     (resumable AI handoff) raises once a long op exhausts its in-process wait budget.
63     """"
64
65     def __init__(self, delay_sec: float, progress: Optional[str] = None) -> None:
66         super().__init__(f"job parked for {delay_sec:.0f}s")
67         self.delay_sec = max(0.0, float(delay_sec))
68         self.progress = progress
69
70
71     def _job_to_request(job: Dict[str, Any]):
72         """Reconstruct the submit-time ProcessRequest from a claimed job row, so a job
73         re-claimed by `claim_due_job` (queued OR resumed-from-waiting) runs identically to
74         the original submit. The row stores exactly the ProcessRequest fields."""
75         import backend.main as _main
76
77         return _main.ProcessRequest(
78             video_path=job["video_path"],
79             player_pool=job.get("player_pool"),
80             max_frames=job.get("max_frames"),
81             segmenter=job.get("segmenter"),
82             # v8 compute-picker: the worker runs _execute_processing ?
71 _maybe_dispatch_remote, which
83             # honours this, so it MUST survive submit?reconstruct (unlike locale, which is
only read off
84             # the job row for analytics). NULL ? server default = pre-picker behaviour.
85             compute=job.get("compute"),
86         )
87
88
89     def _run_claimed_job(job: Dict[str, Any], config: Any, db: Any) -> None:
90         """Run ONE job already claimed (`status='running'`) by `claim_due_job`, recording
its
91         terminal state ? OR parking it for resume on `JobWaiting` (WAIT_AND_RESUME).
92
93         Job `status` is the EXECUTION state (queued|waiting|running|done|failed): 'done'
means
94         the worker finished ? the pipeline's own verdict (validation failure etc.) lives in
95         result["status"]; 'failed' means the worker itself crashed/raised; 'waiting' means
the
96         job self-scheduled and will be re-claimed when `next_poll_at` is due.
97         """"
98         import backend.main as _main
99
100         job_id = job["id"]
101         # Dispatch by kind: 'compile' (#329, the async reel stitch) vs 'process'
(NULL/default ? the
102         # original /api/process pipeline). Both record their terminal state the same way;
only process
103         # jobs feed the M8 cost ledger + M11 flywheel (a compile has no detection
usage/segments).
104         is_compile = job.get("kind") == "compile"
105         try:
106             if is_compile:
107                 result = _main._execute_compile(
108                     config,

```

```

109         db,
110         job,
111         progress=lambda stage: db.set_job_progress(job_id, stage),
112     )
113     else:
114         req = _main._job_to_request(job)
115         result = _main._execute_processing(
116             config,
117             db,
118             req,
119             progress=lambda stage: db.set_job_progress(job_id, stage),
120         )
121         if result.get("video_id"):
122             db.set_job_video_id(job_id, result["video_id"])
123         db.mark_job_done(job_id, result)
124         if not is_compile:
125             _maybe_record_job_cost(
126                 job, result, config, db
127             ) # M8 cost ledger (default-OFF; process jobs only)
128             _maybe_run_flywheel(
129                 job, result, config
130             ) # M11 data-flywheel (process jobs only)
131     except _main.JobWaiting as w:
132         # Cooperative re-enqueue: persist progress + next_poll_at, release the worker.
133         logger.info(f"Job {job_id} parked for {w.delay_sec:.0f}s (resume when due).")
134         db.set_job_waiting(job_id, w.delay_sec, progress=w.progress)
135     except Exception as e:
136         logger.error(f"Job {job_id} crashed: {e}\n{traceback.format_exc()}")
137         db.mark_job_failed(job_id, f"{type(e).__name__}: {e}")
138
139
140     def _maybe_record_job_cost(
141         job: Dict[str, Any], result: Dict[str, Any], config: Any, db: Any
142     ) -> None:
143         """M8 cost ledger (COMPUTE_DECOUPLED_SERVING M8, D8). DEFAULT-OFF: a no-op unless
144         ``config.billing.enabled``. When on, compute the job's cost from ``result['usage']``
145         (= {gpu_seconds, wall_seconds, egress_gb, gemini_usd, gpu_type}) × the versioned
146         pricebook, and
147         persist it.
148
149         ? **Honest gap (M8 scope):** the pipeline does **not** emit ``result['usage']`` yet,
150         so even with
151         ``billing.enabled=True`` AND a real calibrated pricebook, every recorded cost is
152         **0** until a
153         FOLLOW-UP wires the worker to emit usage (gpu/wall/egress/gemini). M8 lands the
154         ledger + pricebook
155         + the recording seam; the usage-emission step is separate (the metering hooks on the
156         real run).
157         With the placeholder pricebook the cost is 0 regardless.
158
159         Best-effort ? never fails a completed job because cost bookkeeping hiccuped. ``job``
160         is the
161         claimed row (carries ``owner_id``); ``result`` is the pipeline output."""
162
163         billing_cfg = getattr(config, "billing", None)
164         if not billing_cfg or not getattr(billing_cfg, "enabled", False):
165             return
166         try:
167             usage = (result or {}).get("usage") or {}
168             db.record_job_cost(
169                 job["id"],
170                 usage=usage,
171                 pricebook_version=billing_cfg.pricebook_version,
172                 gpu_type=usage.get("gpu_type"),
173                 compute_backend=(result or {}).get("compute_backend"),

```

```

168         owner_id=job.get("owner_id"),
169         margin=billing_cfg.margin,
170     )
171     logger.info(
172         "Recorded cost for job %s (pricebook=%s).",
173         job["id"],
174         billing_cfg.pricebook_version,
175     )
176 except Exception as e: # never wedge a completed job on bookkeeping
177     logger.warning(
178         "Cost recording for job %s failed (non-fatal): %s", job.get("id"), e
179     )
180
181
182 def _maybe_run_flywheel(
183     job: Dict[str, Any], result: Dict[str, Any], config: Any
184 ) -> None:
185     """M11 data-flywheel (COMPUTE_DECOUPLED_SERVING M11, D12). DEFAULT-OFF + INERT: a
no-op unless
186     ``config.flywheel.enabled`` AND attested consent ? and consent is a **faked hard-
deny** today
187     (``flywheel.consent_attested`` returns False for every job until counsel + the
consent schema
188     land), so this captures **nothing** in production. When activated it captures the
consented job's
189     artifacts into the per-video workspace (? later labeling ? the golden training set).
190
191     Best-effort ? never fails a completed job because the flywheel hiccuped. ``job`` is
the claimed
192     row (carries ``owner_id``); ``result`` is the pipeline output."""
193
194     flywheel_cfg = getattr(config, "flywheel", None)
195     if not flywheel_cfg or not getattr(flywheel_cfg, "enabled", False):
196         return
197     try:
198         from backend import flywheel
199
200         flywheel.run_for_job(job, result, config)
201     except Exception as e: # never wedge a completed job on the flywheel
202         logger.warning("Flywheel for job %s failed (non-fatal): %s", job.get("id"), e)
203
204
205 def _drain_due_jobs(config: Any, db: Any) -> int:
206     """One worker tick: claim and run every currently-runnable job (queued, or waiting
whose
207     `next_poll_at` is due), earliest-due first, until the table has nothing runnable.
208
209     Each parked job (`JobWaiting` ? status 'waiting') drops out of this drain and is
picked
210     up by a FUTURE drain once due; we schedule that follow-up drain via
`_schedule_drain` so
211     a single-worker box keeps making progress with no central timer (the DB is the
clock).
212
213     Returns the number of jobs that reached a terminal/parked state this tick.
214     """
215     import backend.main as _main
216
217     ran = 0
218     while True:
219         job = db.claim_due_job()
220         if job is None:
221             break
222         ran += 1
223         _main._run_claimed_job(job, config, db)
224     # If anything is parked-but-not-yet-due, arrange a wake-up so it resumes without a

```

```

new
224         # submit. On the single-box executor this is a delayed re-drain; multi-worker just
relies
225         # on the next submit/drain. Best-effort ? never raises into the worker.
226         _main._schedule_next_drain_if_waiting(config, db)
227         return ran
228
229
230     def _schedule_next_drain_if_waiting(config: Any, db: Any) -> None:
231         """If any job is parked in 'waiting', arm a re-drain for when the soonest one is
due, so a
232         self-scheduled job resumes with no further client submit. The drain itself re-checks
233         due-ness, so an early wake is harmless and a clamped wake for a far-future job is
fine."""
234         import backend.main as _main
235
236         try:
237             delay = db.seconds_until_next_due()
238             except Exception as e: # pragma: no cover - defensive; never wedge the worker
239                 logger.error(f"Could not compute next-due delay: {e}")
240                 return
241             if delay is None:
242                 return # nothing waiting
243             # Clamp so a far-future job doesn't pin a thread forever; the drain re-checks due-
ness.
244             _main._arm_wake_timer(max(0.0, min(delay, _main._MAX_DRAIN_WAKE_SEC)), config, db)
245
246
247     def _arm_wake_timer(delay_sec: float, config: Any, db: Any) -> None:
248         """Arm a single daemon timer that re-submits a drain after ``delay_sec`` (no central
249         scheduler ? the DB is the clock). Isolated as the sole timer-creating seam so tests
can
250         monkeypatch it to a no-op and assert the *scheduling decision* without real
waiting."""
251         import backend.main as _main
252
253         timer = threading.Timer(
254             delay_sec,
255             lambda: _main._get_executor().submit(_main._drain_due_jobs, config, db),
256         )
257         timer.daemon = True
258         timer.start()
259
260
261     #: A parked job further out than this is re-checked by a capped wake rather than one
long
262     #: sleeping timer ? keeps the wake cadence bounded without a central scheduler.
263     _MAX_DRAIN_WAKE_SEC = 60.0
264

```

21. user (2026-07-02T04:47:12.227Z)

```

backend\main.py:1106:         "job": os.environ.get("CLOUD_RUN_JOB", "rally-worker"),
backend\worker\_main_.py:209:     scratch = args.scratch or tempfile.mkdtemp(prefix="rally-
worker-")
backend\worker\_main_.py:473:     scratch = args.scratch or tempfile.mkdtemp(prefix="rally-
train-")
backend\compute\__init__.py:66:     job_name = os.environ.get("CLOUD_RUN_JOB", "rally-worker")
backend\compute\cloud_run.py:220:         tempfile.mkdtemp(prefix="cr-read-"), ref.name
or "obj.json"
backend\storage\__init__.py:89:     scratch = tempfile.mkdtemp(prefix="rally-input-",
dir=scratch_parent)
backend\eval\golden_real_fixtures.py:355:         tempfile.mkdtemp(prefix=f".{out_slug}.tmp.",
dir=str(fixtures_dir))

```


22. user (2026-07-02T04:47:20.284Z)

```
180         "returncode": returncode,
181         "segment_count": segment_count,
182         "status": status,
183     }
184     local = os.path.join(scratch, "_done.json")
185     with open(local, "w", encoding="utf-8") as f:
186         json.dump(marker, f)
187     storage.put(local, done_uri, config=storage_cfg)
188     logger.info("[worker] wrote completion marker -> %s", done_uri)
189 except Exception as e: # never fail a good run because the marker push hiccuped
190     logger.warning("[worker] could not write done-marker %s: %s", done_uri, e)
191
192
193 def cmd_infer(args: argparse.Namespace) -> int:
194     config = load_config(args.config) if args.config else load_config()
195     _apply_overrides(config, args.set)
196     if not _run_startup_checks(config):
197         return 2
198
199     # M6: a cloud compute_target DISPATCHES to a Cloud Run Job instead of running in-
process.
200     # DEFAULT-OFF: get_compute_backend returns None for ""/local_gpu/cpu_mock ? fall
through to the
201     # unchanged inline path below. The dispatched container runs THIS cmd_infer inline
(the
202     # dispatcher forces compute_target=local_gpu), so it never recursively re-
dispatches.
203     compute = get_compute_backend(config)
204     if compute is not None:
205         return _dispatch_remote(compute, args)
206
207     storage_cfg = config.storage.model_dump()
208
209     scratch = args.scratch or tempfile.mkdtemp(prefix="rally-worker-")
210     os.makedirs(scratch, exist_ok=True)
211     config.output.output_dir = scratch
212
213     # 1. PULL ? local:// / bare path = identity passthrough; s3:// / gcs:// = download
to scratch.
214     local_video = storage.fetch(
215         args.video_uri, os.path.join(scratch, "in.mp4"), config=storage_cfg
216     )
217     print(f"[worker] pulled {args.video_uri} -> {local_video}")
218
219     # 2. IDENTITY ? whole-file MD5 (the labels?video join key; pulled bytes must be
byte-identical).
220     video_id = compute_video_id(local_video)
221
222     # 3. VALIDATE (same registry as the main CLI).
223     val_results, val_context = validator_registry.run_all_with_context(
224         local_video, config
225     )
226     failures = [r for r in val_results if not r.passed]
227     db = Database(os.path.join(scratch, config.output.db_name))
228     db.add_video(
229         video_id,
230         local_video,
231         "failed" if failures else "processed",
232         [r.model_dump() for r in val_results],
233     )
234
235     def _fail(rc: int) -> int:
236         # M6: on a worker-side FAILURE, write the done-marker (the REAL rc, 0 segments)
```

```

so a remote
237         # dispatcher's bucket-poll terminates FAST + accurately instead of hanging until
its poll
238         # budget expires. DEFAULT-OFF: no --done-uri ? no marker written (the unchanged
local run).
239         if getattr(args, "done_uri", None):
240             _write_done_marker(
241                 args.done_uri, video_id, rc, 0, "failed", storage_cfg, scratch
242             )
243         return rc
244
245     segments: List[dict] = []
246     segmenter_actual = None
247     segmentation_ran = False
248     status = "failed"
249     try:
250         if failures:
251             print(
252                 "[worker] validation FAILED: "
253                 + "; ".join(f"{r.validator_name}: {r.message}" for r in failures)
254             )
255             return _fail(2)
256         seg_name = args.segmenter or config.indexing.default_segmenter
257         segmenter_cls = segmenter_registry.get(seg_name)
258         if not segmenter_cls:
259             print(
260                 f"[worker] unknown segmenter: {seg_name!r} "
261                 f"(have {list(segmenter_registry.list_available())})"
262             )
263             return _fail(2)
264         segmenter_actual = seg_name
265         result = segmenter_cls(db, config).process_video(
266             video_id, local_video, max_frames=args.max_frames
267         )
268         segmentation_ran = True
269         if isinstance(result, Failure):
270             print(f"[worker] segmenter FAILED: {result.message}")
271             return _fail(3)
272         segments = result.value if hasattr(result, "value") else result
273         status = "success"
274         # Loud, never-silent diagnostic: a 0-segment success is almost always a
misconfig, not an
275         # empty match ? the cold-start failure mode (the substrate detector yielded no
usable
276         # trajectories). Explain it so a run is analysable from the log alone (re-run -v
for more).
277         if not segments:
278             detector_impl = (
279                 os.environ.get("RALLY_DETECTOR_IMPL")
280                 or getattr(config.indexing, "detector_impl", None)
281                 or "auto"
282             )
283             logger.warning(
284                 "infer produced 0 segments for video_id=%s (segmenter=%s,
detector_impl=%s). "
285                 "The detector substrate yielded no usable trajectories/windows ? check
the detector "
286                 "logs above; common causes: native runner UNAVAILABLE
(weights/repo/WASB_PYTHON env), "
287                 "the WASB env produced no detections, or the input resolution differs
from the tuned "
288                 "proxy resolution. Re-run with -v for the native command + frame
counts.",
289                 video_id,
290                 segmenter_actual,

```

```

291         detector_impl,
292     )
293     finally:
294         # Best-effort per-video observability report (never raises) ? parity with the
main CLI.
295         emit_indexer_report(
296             db=db,
297             video_id=video_id,
298             video_path=local_video,
299             config=config,

```

23. user (2026-07-02T04:47:43.404Z)

```

backend\reporting\__init__.py:29: def emit_indexer_report(
backend\reporting\__init__.py:74:         output_dir = (config.get("report") or
{ }).get("output_dir")
backend\reporting\__init__.py:75:         res = rec.write(output_dir=output_dir)

```

24. user (2026-07-02T04:47:47.608Z)

```

1     """Per-video observability reports for the indexer.
2
3     `obsreport` is a SOFT dependency: if it is not installed, this module degrades to a
4     no-op and the pipeline runs exactly as before. The single entrypoint
5     `emit_indexer_report(...)` is designed to be called from a `finally` block and
6     NEVER raises ? emitting a report must not be able to break a processing run.
7     """
8
9     from __future__ import annotations
10
11     from typing import Any, Dict, List, Optional
12
13     from . import run_signals # no obsreport dependency ? always importable
14
15     try: # soft dependency ? never hard-fail the pipeline on a reporting import
16         import obsreport # noqa: F401
17
18         from . import harvest
19
20         OBSREPORT_AVAILABLE = True
21     except Exception: # pragma: no cover - exercised only when obsreport is absent
22         OBSREPORT_AVAILABLE = False
23
24
25     def report_enabled(config: Optional[Dict[str, Any]]) -> bool:
26         return bool((config or { }).get("report", { }).get("enabled", True))
27
28
29     def emit_indexer_report(
30         *,
31         db,
32         video_id: str,
33         video_path: Optional[str],
34         config: Optional[Dict[str, Any]],
35         val_results: List[Any],
36         val_context: Optional[Dict[str, Any]] = None,
37         segmenter_requested: Optional[str] = None,
38         segmenter_actual: Optional[str] = None,
39         was_reingest: bool = False,
40         segmentation_ran: bool = True,
41     ) -> Optional[Dict[str, str]]:
42         """Harvest + write the report. Returns the written file paths, or None if
43         reporting is unavailable/disabled/failed. Never raises."""
44         config = config or { }
45         # Accept either a plain dict or a typed config object (e.g.

```

```

backend.config.AppConfig).
46     # Normalize to a plain nested dict so harvest's dict access works uniformly.
47     if hasattr(config, "model_dump"):
48         try:
49             config = config.model_dump(mode="json")
50         except Exception:
51             config = {}
52     # Consume run signals unconditionally (even when reporting is off) so the
53     # process-local store never accumulates entries for finished runs.
54     signals = run_signals.pop(video_id)
55     if not OBSREPORT_AVAILABLE or not report_enabled(config):
56         return None
57     try:
58         play_segments = db.query_play_segments(video_id, {}) if db is not None else []
59         candidates = db.query_candidate_segments(video_id) if db is not None else []
60         rec = harvest.record_run(
61             video_id=video_id,
62             video_path=video_path,
63             config=config,
64             val_results=val_results,
65             val_context=val_context,
66             play_segments=play_segments,
67             candidates=candidates,
68             segmenter_requested=segmenter_requested,
69             segmenter_actual=segmenter_actual,
70             was_reingest=was_reingest,
71             segmentation_ran=segmentation_ran,
72             run_signals=signals,
73         )
74         output_dir = (config.get("report") or {}).get("output_dir")
75         res = rec.write(output_dir=output_dir)
76         return res.paths or None
77     except Exception: # pragma: no cover - defensive; reporting must never break a run
78         return None
79

```

25. assistant (2026-07-02T04:49:11.430Z)

Now let me check the tests that monkeypatch `acquire\_local\_input`, the config models around storage, and `\_maybe\_sweep\_gemini\_orphans` for pattern reference.

26. user (2026-07-02T04:49:13.323Z)

```

80         json.dump(
81             {
82                 "video_id": self.video_id,
83                 "returncode": self.returncode,
84                 "segment_count": len(self.rows),
85                 "status": "success",
86             },
87             f,
88         )
89         return "exec-1"
90
91
92     @pytest.fixture
93     def cloud_env(tmp_path, monkeypatch):
94         """API wired for cloud_run with a LOCAL bucket (the local storage backend makes the
95         bucket read an
96         identity file read ? no GCS). A gs:// input is 'fetched' to a local fake file for
97         hash+validate."""
98         db = Database(str(tmp_path / "t.db")) # noqa: F841
99         monkeypatch.setattr(m, "db_instance", db)
100         monkeypatch.setattr(m, "db_instance", db)
101         bucket = tmp_path / "bucket"

```

```

100     bucket.mkdir()
101     cfg = AppConfig.model_validate(
102         { # noqa: F841
103             "deployment": {"compute_target": "cloud_run"},
104             "storage": {"backend": "local", "output_root": str(bucket)},
105         }
106     )
107     monkeypatch.setattr(m, "global_config", cfg)
108     monkeypatch.setattr(m, "global_config", cfg)
109     monkeypatch.setattr(m, "_get_executor", lambda: _InlineExecutor())
110     # A real gs:// input would be fetched to hash+validate; return a local fake so that
path is offline.
111     fake_local = tmp_path / "fetched.mp4"
112     fake_local.write_bytes(b"FAKEVIDEOBYTES")
113     monkeypatch.setattr(
114         m.storage, "acquire_local_input", lambda path, scfg: str(fake_local)
115     )
116     monkeypatch.setattr(
117         m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {})
118     )
119     return db, bucket, monkeypatch
120
121
122     def _inject_fake(monkeypatch, fake):
123         """Make backend.compute.get_compute_backend (re-imported inside
124         _maybe_dispatch_remote) build a
125         real CloudRunCompute around the FAKE client + the fast no-wait policy + a fixed
token. ``**kw``
126         forwards the per-request ``target_override`` (item 3) so the cloud-picker path
resolves too."""
127         real = compute_pkg.get_compute_backend
128         monkeypatch.setattr(
129             compute_pkg,
130             "get_compute_backend",
131             lambda config, **kw: real(
132                 config, run_client=fake, policy=_FAST, token="tok", **kw
133             ),
134         )
135
136     def _meta(row):
137         md = row.get("metadata")
138         return json.loads(md) if isinstance(md, str) else (md or {})
139
140
141     def test_cloud_dispatch_ingests_segments(cloud_env):
142         db, _bucket, monkeypatch = cloud_env
143         fake = _FakeCloudClient()
144         _inject_fake(monkeypatch, fake)
145
146         out = m._execute_processing(
147             m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
148         )
149
150         assert out["status"] == "success"
151         assert len(out["play_segments"]) == 2
152         rows = db.query_play_segments(out["video_id"], {})
153         assert len(rows) == 2
154         metas = [_meta(r) for r in rows]
155         assert all(md["source"] == "cloud_run" for md in metas)
156         assert {md.get("shot_count") for md in metas} == {5, 3}
157         assert {md.get("ending_reason") for md in metas} == {"winner", "unforced_error"}
158
159

```

27. user (2026-07-02T04:49:14.240Z)

```
350     edge_auth_enabled: bool = False
351     cf_team_domain: str = (
352         "" # e.g. https://<team>.cloudflareaccess.com (issuer + /certs JWKS)
353     )
354     cf_access_aud: str = "" # the CF Access application AUD tag the token must carry
355     # (CORS allow-origins is set via the RALLY_CORS_ALLOW_ORIGINS env var ? read at
import in
356     # backend/main.py so importing the app does no filesystem I/O; default [""].
Credentials stay
357     # OFF ? Cf-Access is a header, not a cookie.)
358
359
360     class StorageConfig(BaseModel):
361         """Pluggable storage backend (docs/DECOUPLED_COMPUTE/). Default ``local`` = today's
362         behaviour (filesystem paths, byte-for-byte). ``s3`` covers AWS + Cloudflare R2 + DO
Spaces
363         + MinIO (via ``endpoint_url``); ``gcs`` = Google Cloud Storage; ``gdrive`` = a
locally MOUNTED
364         Google Drive (``gdrive.mount_root``). Per-backend settings are loose dicts passed to
the constructor ?
365         **secrets are never stored here**, only endpoints / regions / file paths / source-
selectors
366         (boto3/google auth chains supply credentials). See ``backend/storage/``. """
367
368         backend: Literal["local", "s3", "gcs", "gdrive", "gdrive-api"] = "local"
369         # Output root for produced artifacts. "" => fall back to output.output_dir (non-
breaking).
370         output_root: str = ""
371         # --- Input side (COMPUTE_DECOUPLED_SERVING M2; parallel to the output
`backend`/`output_root`).
372         # Default ``local`` / "" keeps today's behaviour: a bare path / ``local://`` URI is
read IN PLACE
373         # (identity, no copy). Set these to read inputs from a bucket / mounted Drive ? a
bare/relative
374         # input name is resolved against ``input_root`` (e.g.
``input_root``="gs://bucket/videos"`` + "x.mp4"
375         # => ``gs://bucket/videos/x.mp4``) or, if ``input_root`` is empty, prefixed with
``input_backend``
376         # (``"gcs"`` + "bucket/x.mp4" => ``gcs://bucket/x.mp4``). An explicit scheme on the
input
377         # (``gs://?``/``s3://?``) or an absolute local path ALWAYS wins. Non-local inputs
are fetched to a
378         # scratch dir before processing (see ``backend.storage.acquire_local_input``).
379         input_backend: Literal["local", "s3", "gcs", "gdrive", "gdrive-api"] = "local"
380         input_root: str = ""
381         # --- Per-video workspace (docs/PER_VIDEO_WORKSPACE.md) ---
382         # One folder per video keyed by the MD5 video_id:
<root>/[<workspace_prefix>]/<video_id>/...
383         # so a local deploy and a cloud (bucket) deploy share ONE relative layout (only the
root
384         # differs). DEFAULT-OFF: convergence is phased + flipped via a Launch Report (D5).
When OFF
385         # the legacy flat layout is used unchanged.
386         single_folder_per_video: bool = False
387         # Durable workspace root URI. "" => precedence: output_root, then output.output_dir.
388         workspace_root: str = ""
389         # Cloud namespace prefix under the root so per-video folders sit tidily beside
390         # videos/ labels/ weights/ (e.g. gs://bucket/workspaces/<video_id>/). "" = root
directly.
391         workspace_prefix: str = "workspaces"
392         local: dict[str, Any] = Field(default_factory=dict)
393         s3: dict[str, Any] = Field(default_factory=dict)
394         gcs: dict[str, Any] = Field(default_factory=dict)
```

```

395         gdrive: dict[str, Any] = Field(default_factory=dict)
396         # gdrive-api settings (e.g. {"root_folder_id": "<id>"}); read via the
hyphen?underscore lookup in
397         # backend.storage.get_storage. NO secrets ? creds come from ADC over the Drive
scope.
398         gdrive_api: dict[str, Any] = Field(default_factory=dict)
399
400
401     class DeploymentConfig(BaseModel):
402         """Deployment profile ? one switch that selects a coherent bundle of compute +
storage
403         knobs (COMPUTE_DECOUPLED_SERVING M1, ``docs/COMPUTE_DECOUPLED_SERVING/README.md``
$4.3).
404
405         **Pure / additive:** the empty default profile (``""``) applies NO preset, so the
pipeline
406         behaves exactly as before this section existed. A profile is opt-in (``config.json``
407         ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); auto-detect only ever
408         *suggests* one (D15 ? cloud spend is never entered silently). Nothing in the core
reads
409         these fields yet (M1): selecting a profile works purely by pre-setting the EXISTING
knobs
410         (``indexing.detector_impl`` / ``skip_ai_handoff``, ``storage.backend``) in the
loader.
411
412         Valid values mirror ``backend/config/presets.py`` (parity-pinned by a unit test)."""
413
414         # "" = no profile (today's manual knobs). Others apply presets.PROFILE_PRESETS.
415         profile: Literal["", "truly-local", "cloud-serving", "cpu-mock"] = ""
416         # Where the core runs. Set by the profile preset; overridable in config.json. Inert
in M1
417         # (recorded for observability; consumed by the ComputeBackend seam from M2+).
418         compute_target: Literal["", "local_gpu", "cloud_run", "cpu_mock"] = ""
419
420
421     class BillingConfig(BaseModel):
422         """Per-job cost ledger (COMPUTE_DECOUPLED_SERVING M8, D8).
423
424         **DEFAULT-OFF:** ``enabled=False`` ? nothing records cost (the v6 cost columns stay
NULL =
425         today's behaviour). When enabled, a job's cost is computed + stored in **USD** via
the versioned
426         ``pricebook`` DB table and **displayed in INR** at ``fx_inr_per_usd``. The default
pricebook
427         version is the seeded ``placeholder-v0`` (every rate ``0.0`` ? cost ``0``) until the
owner
428         inserts a real, calibrated version and points ``pricebook_version`` at it.
429

```

28. user (2026-07-02T04:49:20.813Z)

```

1808: def _maybe_sweep_gemini_orphans(config) -> None:
1809-     """O1/P11: best-effort sweep of Gemini File-API uploads orphaned by a previous crashed
run.
1810-     A no-op unless the configured provider is Gemini AND its API key is present. Never
raises ?
1811-     a failed sweep is logged and swallowed. The Gemini client is built with an HTTP timeout
so
1812-     a hung connection can't stall startup indefinitely. Extracted for unit tests."""
1813-     try:
1814-         provider_name = getattr(config, "ai_provider", "")
1815-         if provider_name != "gemini":
1816-             return
1817-         api_key = os.environ.get("GEMINI_API_KEY")
1818-         if not api_key:

```

```

1819-         return
1820-     provider = provider_registry.get(
1821-         "gemini", api_key=api_key, model_name=config.ai_model
1822-     )
1823-     deleted = provider.cleanup_orphaned_files()
1824-     if deleted:
1825-         logger.info(
1826-             "[STARTUP] Gemini orphan sweep reclaimed %d remote file(s).", deleted
1827-         )
1828- except Exception as e:
1829-     logger.warning("[STARTUP] Gemini orphan sweep skipped (non-fatal): %s", e)
1830-
1831-
1832-def main():
1833-     parser = argparse.ArgumentParser(
1834-         description="Sports Highlight Indexer Orchestrator"
1835-     )
1836-     parser.add_argument("--video-path", help="Local path to the MP4 sports video file")
1837-     parser.add_argument(
1838-         "--output-dir",
1839-         default="output",
1840-         help="Directory where processed clips/highlights are saved",
1841-     )
1842-     parser.add_argument(
1843-         "--setup",
1844-         action="store_true",
1845-         help="Launch dependency checker and setup utility",
1846-     )
1847-     parser.add_argument(
1848-         "--debug-clip",

```

29. user (2026-07-02T04:49:27.020Z)

```

1     """Input acquisition for COMPUTE_DECOUPLED_SERVING M2 ? resolve_input_uri /
input_is_local /
2     acquire_local_input.
3
4     The hard guarantee: the default config (input_backend='local', input_root='') is a byte-
for-byte
5     identity passthrough for every input path used today (no fetch, no copy). A bucket/Drive
source is
6     routed through storage.fetch to a local scratch file. All offline (the remote backend is
a
7     monkeypatched fetch ? no boto3/google-cloud-storage, no network).
8
9     OS note: os.path.isabs is platform-dependent, so the absolute-path cases use a POSIX-
style
10    '/data/...' (absolute on BOTH Windows and POSIX) and Windows backslash paths only with
input_root=''
11    (branch 5, OS-independent) ? the full suite runs cross-OS in CI.
12    """
13
14    import os
15
16    import pytest
17
18    from backend import storage
19    from backend.storage import acquire_local_input, input_is_local, resolve_input_ref
20    from backend.storage.ref import resolve_input_uri
21
22
23    # ----- resolve_input_uri
24    @pytest.mark.parametrize(
25        "raw,ib,ir,expected",
26        [

```



```

27         # 1. explicit scheme always wins (even over a non-local backend + root)
28         ("gs://b/x.mp4", "local", "", "gs://b/x.mp4"),
29         ("s3://b/x.mp4", "gcs", "gs://r", "s3://b/x.mp4"),
30         ("local://abs/x.mp4", "gcs", "gs://r", "local://abs/x.mp4"),
31         ("gdrive://Sharing/x.mp4", "local", "", "gdrive://Sharing/x.mp4"),
32         # 2. an anchored local path (leading slash, or drive prefix) is always local,
never joined ?
33         # OS-independently (not via os.path.isabs, which differs on Windows+py3.13)
34         ("/data/x.mp4", "gcs", "gs://r", "/data/x.mp4"),
35         ("\\\\server\\share\\x.mp4", "gcs", "gs://r", "\\\\server\\share\\x.mp4"),
36         ("C:\\v\\x.mp4", "gcs", "gs://r", "C:\\v\\x.mp4"),
37         ("C:/v/x.mp4", "gcs", "gs://r", "C:/v/x.mp4"),
38         # 3. a bare/relative name joins under a non-empty input_root
39         ("x.mp4", "local", "gs://bucket/videos", "gs://bucket/videos/x.mp4"),
40         (
41             "sub/x.mp4",
42             "local",
43             "gs://bucket/videos/",
44             "gs://bucket/videos/sub/x.mp4",
45         ), # root trailing slash
46         # 4. input_backend prefixes a bare bucket/key when no input_root
47         ("bucket/x.mp4", "gcs", "", "gs://bucket/x.mp4"),
48         ("bucket/x.mp4", "s3", "", "s3://bucket/x.mp4"),
49         # 5. bare/relative local path ? today's behaviour, unchanged
50         ("x.mp4", "local", "", "x.mp4"),
51         ("rel/sub/x.mp4", "local", "", "rel/sub/x.mp4"),
52         (
53             "C:\\v\\x.mp4",
54             "local",
55             "",
56             "C:\\v\\x.mp4",
57         ), # Windows path, default config: unchanged (OS-independent)
58     ],
59 )
60 def test_resolve_input_uri(raw, ib, ir, expected):
61     assert resolve_input_uri(raw, input_backend=ib, input_root=ir) == expected
62
63
64 def test_resolve_input_uri_default_is_identity():
65     for p in ("/data/clip.mp4", "rel/clip.mp4", "clip.mp4", "gs://b/k", "local://x"):
66         assert resolve_input_uri(p) == p
67
68
69 # ----- input_is_local
70 def test_input_is_local_true_for_local_paths():
71     assert input_is_local("/data/x.mp4") is True
72     assert input_is_local("rel/x.mp4") is True
73     assert input_is_local("local://x.mp4") is True
74
75
76 def test_input_is_local_false_for_buckets():
77     assert input_is_local("gs://b/x.mp4") is False
78     assert input_is_local("s3://b/x.mp4") is False
79     assert input_is_local("gdrive://Sharing/x.mp4") is False
80
81
82 def test_input_is_local_follows_input_root():
83     cfg = {"input_backend": "local", "input_root": "gs://bucket/videos"}
84     assert input_is_local("x.mp4", cfg) is False # bare name resolves into the bucket
85     assert input_is_local("/abs/x.mp4", cfg) is True # absolute path stays local
86     assert input_is_local("gs://other/x.mp4", cfg) is False
87
88
89 def test_input_is_local_accepts_pydantic_storage_config():
90     from backend.config.models import StorageConfig

```

```

91
92     sc = StorageConfig(input_root="gs://bucket/videos")
93     assert input_is_local("x.mp4", sc) is False
94     assert input_is_local("/abs/x.mp4", sc) is True
95
96
97     def test_input_is_local_unknown_scheme_raises():
98         # An unsupported scheme is a ValueError (callers at the API/CLI boundary map it to a
99         # never a silent True/False that could mislead a security gate.
100         for bad in ("https://h/x.mp4", "ftp://h/x.mp4", "file:///etc/passwd"):
101             with pytest.raises(ValueError):
102                 input_is_local(bad)
103
104
105     # ----- resolve_input_ref
106     def test_resolve_input_ref_local_and_remote():
107         local = resolve_input_ref("/data/x.mp4")
108         assert (local.backend, local.key) == ("local", "/data/x.mp4")
109         # local:// strips to the bare key (so callers stat ref.key, not the raw URI)
110         assert resolve_input_ref("local:///data/x.mp4").key == "/data/x.mp4"
111         remote = resolve_input_ref("x.mp4", {"input_root": "gs://bucket/videos"})
112         assert (remote.backend, remote.bucket, remote.key) == (
113             "gcs",
114             "bucket",
115             "videos/x.mp4",
116         )
117
118
119     def test_resolve_input_ref_unknown_scheme_raises():
120         with pytest.raises(ValueError):
121             resolve_input_ref("https://host/x.mp4")
122
123
124     # ----- acquire_local_input
125     def test_acquire_local_input_local_identity(tmp_path):
126         p = tmp_path / "v.mp4"
127         p.write_bytes(b"abc")
128         # Identity passthrough ? returns the SAME string, no copy, no temp dir.
129         assert acquire_local_input(str(p)) == str(p)
130
131
132     def test_acquire_local_input_relative_local_identity():
133         assert acquire_local_input("rel/clip.mp4") == "rel/clip.mp4"
134
135
136     def _capturing_fetch(captured):
137         def _f(uri, dst=None, *, config=None, overwrite=False):
138             captured["uri"] = uri
139             captured["dst"] = dst
140             captured["config"] = config
141             if dst:
142                 with open(dst, "wb") as fh:
143                     fh.write(b"remote-bytes")
144             return dst or uri
145
146         return _f
147
148
149     def test_acquire_local_input_remote_fetches_to_scratch(monkeypatch):
150         captured = {}
151         monkeypatch.setattr(storage, "fetch", _capturing_fetch(captured))
152         out = acquire_local_input(
153             "gs://bucket/videos/x.mp4", {"backend": "local", "gcs": {}}
154         )

```

```

155     assert captured["uri"] == "gs://bucket/videos/x.mp4" # exact URI routed to fetch
156     assert os.path.basename(out) == "x.mp4" # named from the ref
157     assert os.path.isfile(out) and open(out, "rb").read() == b"remote-bytes"
158     # The whole storage settings dict is forwarded so per-backend (gcs/s3) config
reaches fetch.
159     assert captured["config"] == {"backend": "local", "gcs": {}}
160
161
162     def test_acquire_local_input_input_root_joins_then_fetches(monkeypatch):
163         captured = {}
164         monkeypatch.setattr(storage, "fetch", _capturing_fetch(captured))
165         acquire_local_input("x.mp4", {"input_root": "gs://bucket/videos"})
166         assert captured["uri"] == "gs://bucket/videos/x.mp4" # resolved BEFORE fetch
167
168
169     def test_acquire_local_input_never_fetches_for_local(monkeypatch):
170         def _boom(*a, **k):
171             raise AssertionError("fetch must not be called for a local input")
172
173         monkeypatch.setattr(storage, "fetch", _boom)
174         assert acquire_local_input("/data/x.mp4") == "/data/x.mp4"
175         assert (
176             acquire_local_input("clip.mp4", {"input_backend": "local", "input_root": ""})
177             == "clip.mp4"
178         )
179

```

30. user (2026-07-02T04:50:56.953Z)

```

360         ) # explicit "run on my machine" ? no dispatch despite a cloud server
361
362
363     def test_per_request_cloud_forces_dispatch_on_local_server(tmp_path, monkeypatch):
364         """Server default is LOCAL (compute_target=''), but a per-request compute='cloud'
forces a Cloud
365         Run dispatch when the control plane is wired (CLOUD_RUN_JOB + GOOGLE_CLOUD_PROJECT +
a bucket)."""
366         db = Database(str(tmp_path / "t.db")) # noqa: F841
367         monkeypatch.setattr(m, "db_instance", db)
368         bucket = tmp_path / "bucket"
369         bucket.mkdir()
370         cfg = AppConfig.model_validate(
371             {
372                 # noqa: F841
373                 "deployment": {"compute_target": ""}, # server default = in-process
374                 "storage": {"backend": "local", "output_root": str(bucket)},
375             }
376         )
377         monkeypatch.setattr(m, "global_config", cfg)
378         monkeypatch.setattr(m, "_get_executor", lambda: _InlineExecutor())
379         fake_local = tmp_path / "fetched.mp4"
380         fake_local.write_bytes(b"FAKEVIDEOBYTES")
381         monkeypatch.setattr(
382             m.storage, "acquire_local_input", lambda path, scfg: str(fake_local)
383         )
384         monkeypatch.setattr(
385             m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {})
386         )
387         monkeypatch.setenv("CLOUD_RUN_JOB", "rally-worker")
388         monkeypatch.setenv("GOOGLE_CLOUD_PROJECT", "proj")
389         fake = _FakeCloudClient()
390         _inject_fake(monkeypatch, fake)
391
392     out = m._execute_processing(
393         m.global_config,
394         m.db_instance,

```

```

394         m.ProcessRequest(
395             video_path="gs://b/clip.mp4", segmenter="wasb_hybrid", compute="cloud"
396         ),
397     )
398     assert out["status"] == "success"
399     assert len(out["play_segments"]) == 2
400     assert (
401         len(fake.calls) == 1
402     ) # per-request override dispatched despite the LOCAL server default
403
404
405     def test_per_request_cloud_unconfigured_fails_clearly(tmp_path, monkeypatch):
406         """compute='cloud' on a server NOT wired for it fails clearly (no silent local
407         fallback, no
408         dispatch) so the picker can surface 'cloud isn't available here'."""
409         db = Database(str(tmp_path / "t.db")) # noqa: F841
410         monkeypatch.setattr(m, "db_instance", db)

```

31. user (2026-07-02T04:50:57.684Z)

```

535     ) # in-process rows, not cloud
536
537
538     def test_cloud_override_survives_api_boundary_on_local_server(tmp_path, monkeypatch):
539         """REGRESSION (mirror): on a LOCAL-default server wired for cloud (env + bucket),
540         POST /api/process
541         with compute='cloud' must survive the boundary and dispatch to Cloud Run."""
542         db = Database(str(tmp_path / "t.db")) # noqa: F841
543         monkeypatch.setattr(m, "db_instance", db)
544         bucket = tmp_path / "bucket"
545         bucket.mkdir()
546         cfg = AppConfig.model_validate(
547             { # noqa: F841
548                 "deployment": {"compute_target": ""}, # server default = in-process
549                 "storage": {"backend": "local", "output_root": str(bucket)},
550             }
551         )
552         monkeypatch.setattr(m, "global_config", cfg)
553         monkeypatch.setattr(m, "_get_executor", lambda: _InlineExecutor())
554         fake_local = tmp_path / "fetched.mp4"
555         fake_local.write_bytes(b"FAKEVIDEOBYTES")
556         monkeypatch.setattr(
557             m.storage, "acquire_local_input", lambda path, scfg: str(fake_local)
558         )
559         monkeypatch.setattr(
560             m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {})
561         )
562         monkeypatch.setenv("CLOUD_RUN_JOB", "rally-worker")
563         monkeypatch.setenv("GOOGLE_CLOUD_PROJECT", "proj")
564         fake = _FakeCloudClient(video_id="vidBoundary")
565         _inject_fake(monkeypatch, fake)
566         app = m.create_app(m.global_config, db)
567         client = TestClient(app)
568
569         r = client.post(
570             "/api/process",
571             json={
572                 "video_path": "gs://b/clip.mp4",
573                 "segmenter": "wasb_hybrid",
574                 "compute": "cloud",
575             },
576         )
577         assert r.status_code == 202
578         job = client.get(f"/api/jobs/{r.json()['job_id']}").json()
579         assert job["status"] == "done"

```

```

579         assert (
580             len(fake.calls) == 1
581         ) # 'cloud' honoured across the boundary ? dispatched despite local default
582         segs = db.query_play_segments(job["video_id"], {})
583         assert len(segs) == 2
584         assert all(_meta(s)["source"] == "cloud_run" for s in segs)
585
586
587     # --- compute-path PROVENANCE (the validation method: prove which backend ran)
588     -----
589
590     def test_cloud_dispatch_stamps_compute_provenance(cloud_env):
591         """A cloud run stamps result['compute'] with path=cloud_run + the REAL Cloud Run
592         execution name
593         (the GCP-verifiable proof) + the requested choice + job/region/outputs ? so the
594         intended path is
595         provable from the job result alone (no guessing from segment metadata)."""
596         _db, _bucket, monkeypatch = cloud_env

```

32. user (2026-07-02T04:51:05.455Z)

```

tests\test_async_jobs.py:391: def test_worker_dispatches_compile_kind_to_execute_compile(env):
tests\test_async_jobs.py:392:     """The claimed-job dispatcher routes kind='compile' to
tests\test_async_jobs.py:405:     _execute_compile (NOT _execute_processing),
tests\test_api_assets.py:1: """TestClient coverage for POST /api/assets ? "store this video in my
medium" (STUDIO_MEDIA_LIFECYCLE.md P4, D14).
tests\test_api_assets.py:46:         "/api/assets", json={"uploaded_path": str(up), "medium_uri":
"local://"}
tests\test_api_assets.py:72:         "/api/assets", json={"uploaded_path": str(up), "medium_uri":
scheme}
tests\test_api_assets.py:96:         "/api/assets", json={"uploaded_path": str(up), "medium_uri":
"gdrive://videos"}
tests\test_api_assets.py:114:         "/api/assets",
tests\test_api_assets.py:131:         "/api/assets", json={"uploaded_path": str(up),
"medium_uri": "s3://b/k"}
tests\test_api_assets.py:141:         "/api/assets",
tests\test_api_assets.py:144:         neither = client.post("/api/assets", json={"medium_uri":
"s3://b/k"})
tests\test_api_assets.py:151:         "/api/assets",
tests\test_api_assets.py:160:         r = client.post("/api/assets", json={"uploaded_path": str(up),
"medium_uri": ""})
tests\test_api_assets.py:180:         "/api/assets", json={"uploaded_path": str(up),
"medium_uri": "gdrive://videos"}
tests\test_api_assets.py:194:         "/api/assets", json={"uploaded_path": str(up),
"medium_uri": "s3://b/k"}
tests\test_api_assets.py:204:         "/api/assets", json={"uploaded_path": str(up),
"medium_uri": "ftp://nope/x"}
tests\test_compute_backend.py:202: def test_read_json_passes_dst_for_cloud_backends(monkeypatch):
tests\test_compute_backend.py:203:     """BLOCKER guard: GCS/S3 fetch_to_local REQUIRES a dst (a
None dst raises) ? _read_json must
tests\test_compute_backend.py:222:     data = backend._read_json("gs://b/x.json")
tests\test_golden_real_fixtures.py:157:     prov = gf._read_json(fixtures / slug /
"provenance.json")
tests\test_golden_real_fixtures.py:291:     prov = gf._read_json(fixtures / expected["slug"] /
"provenance.json")
tests\test_golden_real_fixtures.py:319:     prov = gf._read_json(fixtures / ok["slug"] /
"provenance.json")
tests\test_golden_real_fixtures.py:541:     committed = gf._read_json(fdir / "expected.json")
tests\test_golden_real_fixtures.py:746:     prov = gf._read_json(fixtures / slug /
"provenance.json")
tests\test_golden_real_fixtures.py:794:     gf._read_json(fixtures / "fx_trim01" /
"expected.json"),

```

33. user (2026-07-02T04:51:12.900Z)

```
190     backend = CloudRunCompute(run_client=client, job_name="j", region="r", policy=_FAST)
191     backend.run_infer(
192         ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bk"))
193     )
194     backend.run_infer(
195         ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bk"))
196     )
197     done1 = client.calls[0]["argv"][client.calls[0]["argv"].index("--done-uri") + 1]
198     done2 = client.calls[1]["argv"][client.calls[1]["argv"].index("--done-uri") + 1]
199     assert done1 != done2
200
201
202     def test_read_json_passes_dst_for_cloud_backends(monkeypatch):
203         """BLOCKER guard: GCS/S3 fetch_to_local REQUIRES a dst (a None dst raises) ?
_read_json must
204         pass one, else every real cloud marker/report read silently returns {} (local hid
this via its
205         identity passthrough). Mimic a cloud backend that rejects dst=None."""
206         from backend.compute import cloud_run as cr
207
208         captured = {}
209
210         def fake_fetch(uri, dst=None, *, config=None, overwrite=False):
211             captured["dst"] = dst
212             if dst is None:
213                 raise ValueError(
214                     "cloud backend requires a dst path"
215                 ) # mirror GCSStorage/S3Storage
216             with open(dst, "w", encoding="utf-8") as f:
217                 f.write('{"video_id": "v", "returncode": 0}')
218             return dst
219
220         monkeypatch.setattr(cr, "fetch", fake_fetch)
221         backend = cr.CloudRunCompute(run_client=_FakeRunClient(), job_name="j", region="r")
222         data = backend._read_json("gs://b/x.json")
223         assert (
224             captured["dst"] is not None
225         ) # the fix: a real dst is always passed (no silent {})
226         assert data == {"video_id": "v", "returncode": 0}
227
228
229     def test_cloud_run_polls_until_marker_appears(monkeypatch):
230         """The bucket-poll WAIT path (D7 scale-to-zero cold start): the marker is ABSENT for
the first
231         polls then appears ? poll_until must LOOP, not resolve on poll 0. The synchronous
fake elsewhere
232         never exercises this; here exists() is False twice then True."""
233         from backend.compute import cloud_run as cr
234
235         exists_calls = {"n": 0}
236
237         class _FakeStore:
238             def exists(self, ref):
239                 exists_calls["n"] += 1
```

34. user (2026-07-02T04:51:13.996Z)

```
100     assert body["medium"] == "gdrive" and body["copied"] is True
101     assert body["stored_uri"] == f"gdrive://videos/{vid}.mp4"
102     # the real copy landed on the mount filesystem
103     assert (mount / "videos" / f"{vid}.mp4").read_bytes() == up.read_bytes()
104
105
```

```

106 def test_source_uri_already_at_rest_registers_without_copy(tmp_path, monkeypatch):
107     client, db, _ = _client(tmp_path, monkeypatch)
108     at_rest = _upload(tmp_path, name="already.mp4") # a local path already at rest
109     puts = []
110     monkeypatch.setattr(
111         m.storage, "put", lambda *a, **k: puts.append(a) or "should-not-be-called"
112     )
113     r = client.post(
114         "/api/assets",
115         json={"source_uri": str(at_rest), "medium_uri": "local://"},
116     )
117     assert r.status_code == 201
118     body = r.json()
119     assert body["copied"] is False and puts == [] # no copy (D14b)
120     assert body["stored_uri"] == str(at_rest)
121     assert body["video_id"] == compute_video_id(str(at_rest))
122
123
124 # ---- guards -----
125
126
127 def test_self_host_posture_rejects_when_jailed(tmp_path, monkeypatch):
128     client, db, _ = _client(tmp_path, monkeypatch, allowed_video_root=str(tmp_path))
129     up = _upload(tmp_path)
130     r = client.post(
131         "/api/assets", json={"uploaded_path": str(up), "medium_uri": "s3://b/k"}
132     )
133     assert r.status_code == 403
134     assert "hosted" in r.json()["detail"].lower()
135
136
137 def test_requires_exactly_one_source(tmp_path, monkeypatch):
138     client, db, _ = _client(tmp_path, monkeypatch)
139     up = _upload(tmp_path)
140     both = client.post(
141         "/api/assets",
142         json={"uploaded_path": str(up), "source_uri": "s3://b/k", "medium_uri":
"s3://b/k"},
143     )
144     neither = client.post("/api/assets", json={"medium_uri": "s3://b/k"})
145     assert both.status_code == 400 and neither.status_code == 400
146
147
148 def test_missing_uploaded_file_404(tmp_path, monkeypatch):
149     client, db, _ = _client(tmp_path, monkeypatch)
150     r = client.post(
151         "/api/assets",
152         json={"uploaded_path": str(tmp_path / "nope.mp4"), "medium_uri": "s3://b/k"},
153     )
154     assert r.status_code == 404
155
156
157 def test_upload_requires_medium_uri(tmp_path, monkeypatch):
158     client, db, _ = _client(tmp_path, monkeypatch)
159     up = _upload(tmp_path)
160     r = client.post("/api/assets", json={"uploaded_path": str(up), "medium_uri": ""})
161     assert r.status_code == 400
162     assert "medium_uri" in r.json()["detail"]
163
164
165 def test_free_space_507_on_low_disk_drive_medium(tmp_path, monkeypatch):
166     mount = tmp_path / "drive"
167     mount.mkdir()
168     db = Database(str(tmp_path / "t.db"))
169     cfg = AppConfig()

```

35. user (2026-07-02T04:51:51.412Z)

```
tests\test_api_endpoints.py:234: mp.setattr(m.segmenter_registry, "get", lambda name:
_FakeSeg)
tests\test_api_endpoints.py:250: mp.setattr(m.segmenter_registry, "get", lambda name:
_FakeSeg)
tests\test_api_endpoints.py:312: mp.setattr(m.segmenter_registry, "get", lambda name:
_FakeSeg)
tests\test_api_endpoints.py:660: mp.setattr(m.segmenter_registry, "get", lambda name:
_FakeSeg)
tests\test_api_endpoints.py:711: mp.setattr(m.segmenter_registry, "get", lambda name:
_CaptureSeg)
tests\test_api_endpoints.py:738: mp.setattr(m.segmenter_registry, "get", lambda name:
_CaptureSeg)
tests\test_api_endpoints.py:792: mp.setattr(m.segmenter_registry, "get", lambda name:
_CaptureSeg)
tests\test_api_endpoints.py:823: m.segmenter_registry, "get", lambda name: _FakeSeg
tests\test_api_endpoints.py:866: mp.setattr(m.segmenter_registry, "get", lambda name:
_CaptureSeg)
tests\test_api_endpoints.py:893: mp.setattr(m.segmenter_registry, "get", lambda name:
_FakeSeg)
tests\test_api_endpoints.py:917: mp.setattr(m.segmenter_registry, "get", lambda name:
_FakeSeg)
tests\test_api_endpoints.py:937: monkeypatch.setattr(m.segmenter_registry, "get", lambda
name: _CaptureSeg)
tests\test_api_endpoints.py:989: monkeypatch.setattr(m.segmenter_registry, "get", lambda
name: _NoSeg)
tests\test_api_cloud_dispatch.py:263: monkeypatch.setattr(m.segmenter_registry, "get", lambda
name: _InProcSeg)
tests\test_api_cloud_dispatch.py:346: monkeypatch.setattr(m.segmenter_registry, "get", lambda
name: _InProcSeg)
tests\test_api_cloud_dispatch.py:513: monkeypatch.setattr(m.segmenter_registry, "get", lambda
name: _InProcSeg)
tests\test_api_cloud_dispatch.py:617: monkeypatch.setattr(m.segmenter_registry, "get", lambda
name: _InProcSeg)
tests\test_api_cloud_dispatch.py:639: monkeypatch.setattr(m.segmenter_registry, "get", lambda
name: _InProcSeg)
tests\test_async_jobs.py:168: mp.setattr(m.segmenter_registry, "get", lambda name: _FakeSeg)
tests\test_async_jobs.py:223: mp.setattr(m.segmenter_registry, "get", lambda name: _EmptySeg)
tests\test_async_jobs.py:241: mp.setattr(m.segmenter_registry, "get", lambda name: _FakeSeg)
tests\test_async_jobs.py:267: mp.setattr(m.segmenter_registry, "get", lambda name: None)
tests\test_async_jobs.py:268: mp.setattr(m.segmenter_registry, "list_available", lambda:
{"motion": "..."})
tests\test_async_jobs.py:285: mp.setattr(m.segmenter_registry, "get", lambda name: _FakeSeg)
tests\test_async_jobs.py:332: mp.setattr(m.segmenter_registry, "get", lambda name:
_RaisingSeg)
tests\test_async_jobs.py:350: mp.setattr(m.segmenter_registry, "get", lambda name: _FakeSeg)
tests\test_async_jobs.py:458: mp.setattr(m.segmenter_registry, "get", lambda name: _FakeSeg)
tests\test_base.py:7: class DummySegmenter(BasePlaySegmenter):
tests\test_base.py:65: class ComplexCtorSegmenter(BasePlaySegmenter):
tests\test_base.py:105: class BadDescSegmenter(BasePlaySegmenter):
tests\test_base.py:191: class RaisingSegmenter(BasePlaySegmenter):
tests\test_generate_highlights.py:84: ), patch.object(gh, "segmenter_registry",
seg_registry):
tests\test_generate_highlights.py:126: gh, "segmenter_registry", seg_registry
tests\test_generate_highlights.py:163: ), patch.object(gh, "segmenter_registry",
seg_registry):
tests\test_generate_highlights.py:260: gh, "segmenter_registry", seg_registry
tests\test_generate_highlights.py:304: gh, "segmenter_registry", seg_registry
tests\test_job_worker_loop.py:85: mp.setattr(m.segmenter_registry, "get", lambda name:
_FakeSeg)
tests\test_job_worker_loop.py:163: mp.setattr(m.segmenter_registry, "get", lambda name:
_WaitOnceSeg)
tests\test_job_worker_loop.py:274: mp.setattr(m.segmenter_registry, "get", lambda name:
_BoomSeg)
```



```
tests\test_light_tier_no_torch.py:37:from backend.pipeline.segmenters.base import
segmenter_registry
```

[Showing results with pagination = limit: 40]

36. user (2026-07-02T04:52:04.559Z)

```
100         assert {
101             "id",
102             "video_path",
103             "video_id",
104             "status",
105             "progress",
106             "error",
107             "result",
108         } <= cols
109         assert {"policy", "next_poll_at"} <= cols # v3 self-scheduling columns
110         assert {
111             "locale",
112             "compute",
113             "kind",
114             "params",
115         } <= cols # v7 locale, v8 compute, v9 compile job
116
117
118     def test_migration_upgrades_v1_db_in_place(tmp_path):
119         path = str(tmp_path / "v1.db")
120         conn = sqlite3.connect(path)
121         conn.execute(
122             "CREATE TABLE videos (id TEXT PRIMARY KEY, filepath TEXT NOT NULL, "
123             "processed_at TIMESTAMP, status TEXT NOT NULL, validation_results TEXT)"
124         )
125         # A genuine v1 DB also has candidate_segments (created in the version<1 block);
include it
126         # so the v4 personalization ALTER (which adds user_id to that table) has a table to
alter.
127         conn.execute(
128             "CREATE TABLE candidate_segments (id INTEGER PRIMARY KEY AUTOINCREMENT, "
129             "video_id TEXT NOT NULL, detector TEXT NOT NULL, start_time REAL NOT NULL, "
130             "end_time REAL NOT NULL, duration REAL NOT NULL, stats TEXT)"
131         )
132         conn.execute("PRAGMA user_version = 1")
133         conn.commit()
134         conn.close()
135
136         db = Database(path) # ladder should take it 1 ? 9
137         with db._get_connection() as c:
138             assert c.execute("PRAGMA user_version").fetchone()[0] == 9
139             assert c.execute(
140                 "SELECT name FROM sqlite_master WHERE type='table' AND name='jobs'"
141             ).fetchone()
142             # Restore v5 personalization coverage (user_models table + user_id on
candidate_segments).
143             # These were dropped in prior edit; add (do not replace) per review.
144             tables = {
145                 r[0]
146                 for r in c.execute(
147                     "SELECT name FROM sqlite_master WHERE type='table'"
148                 ).fetchall()
149             }
150             assert "user_models" in tables
151             cand_cols = {
152                 r["name"]
153                 for r in c.execute("PRAGMA table_info(candidate_segments)").fetchall()
154             }
```

```

155         assert "user_id" in cand_cols
156
157
158     def test_parity_job_model_with_direct_processing(env, tmp_path, monkeypatch):
159         """Item 5 positive: direct CLI-style path (m._execute_processing) vs job-worker path
160         (submit + _InlineExecutor leading to _run_claimed_job / drain, plus explicit
161         _run_claimed_job)
162         produce equivalent outcomes for the same ProcessRequest shape. Exercises success,
163         Failure, and empty-run branches. Real DB queries using vid (no self-compares or
164         tautologies).
165         """
166         client, db, tpath, mp = env
167         vid_file = _video(tpath)
168         vid = m.compute_video_id(str(vid_file))
169         mp.setattr(m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
170         mp.setattr(m.segmenter_registry, "get", lambda name: _FakeSeg)
171
172         req = m.ProcessRequest(video_path=str(vid_file), segmenter="x")
173
174         # Direct CLI-style path
175         direct_out = m._execute_processing(m.global_config, m.db_instance, req)
176         segs = db.query_play_segments(vid, {})
177         assert len(segs) >= 1, f"direct path must affect segments for vid {vid}"
178         assert direct_out.get("result", {}).get("status") in (
179             None,
180             "success",
181             "processed",
182         ) or "play_segments" in str(direct_out)
183
184         # Job-worker path via public API submit (exercises job creation + inline executor +
185         internal _run_claimed_job)
186         rj = client.post(
187             "/api/process", json={"video_path": str(vid_file), "segmenter": "x"}
188         )
189         job_out = client.get(f"/api/jobs/{rj.json()['job_id']}").json()
190         segs_j = db.query_play_segments(vid, {})
191         assert job_out["result"]["status"] == "success"
192         assert len(segs_j) >= 1
193
194         # Explicit job-worker via _run_claimed_job (create + mark running + drive the
195         claimed runner)
196         vid_file2 = _video(tpath, name="m2.mp4")
197         vid2 = m.compute_video_id(str(vid_file2))
198         jid = f"parity-{vid2[:8]}"
199         db.create_job(jid, str(vid_file2))
200         db.mark_job_running(jid)
201         job_row = db.get_job(jid)
202         m._run_claimed_job(job_row, m.global_config, m.db_instance)
203         segs2 = db.query_play_segments(vid2, {})
204         assert len(segs2) >= 1, (
205             f"_run_claimed_job path must produce segments for vid {vid2}"
206         )
207
208         # Failure branch (direct + job submit)
209         mp.setattr(m.registry, "run_all_with_context", lambda *a, **k: (_fail(), {}))
210         req_f = m.ProcessRequest(video_path=str(vid_file))
211         direct_f = m._execute_processing(m.global_config, m.db_instance, req_f)
212         rf = client.post("/api/process", json={"video_path": str(vid_file)})
213         jobf = client.get(f"/api/jobs/{rf.json()['job_id']}").json()
214         assert (
215             direct_f.get("result", {}).get("status") == "failed"
216             or jobf["result"]["status"] == "failed"
217         )
218
219         # Empty-run branch (success verdict but zero play_segments from this execution)

```

```

216 class _EmptySeg:
217     def __init__(self, db, cfg):
218         self.db = db
219
220     def process_video(self, video_id, video_path, *a, **k):
221         return []
222
223 mp.setattr(m.segmenter_registry, "get", lambda name: _EmptySeg)
224 mp.setattr(m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
225 req_e = m.ProcessRequest(video_path=str(vid_file))
226 de = m._execute_processing(m.global_config, m.db_instance, req_e)
227 re = client.post("/api/process", json={"video_path": str(vid_file)})
228 je = client.get(f"/api/jobs/{re.json()['job_id']}").json()
229 assert de.get("result", {}).get("status") in (None, "success")
230 assert je["result"]["status"] == "success"
231 # (reingest safety for prior on empty/fail covered in dedicated tests)
232
233
234 # --- submit-time fast-fail -----
235
236
237 def test_submit_returns_202_with_job_id(env):
238     client, db, tmp_path, mp = env
239     vid_file = _video(tmp_path)
240     mp.setattr(m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
241     mp.setattr(m.segmenter_registry, "get", lambda name: _FakeSeg)
242     r = client.post(
243         "/api/process", json={"video_path": str(vid_file), "segmenter": "x"}
244     )
245     assert r.status_code == 202
246     body = r.json()
247     assert body["job_id"]
248     assert body["status"] == "queued"
249     assert body["poll"].endswith(body["job_id"])

```

37. user (2026-07-02T04:52:10.866Z)

```

1     """Async job model (DEFERRED_HARDENING_PLAN #16).
2
3     Contract under test:
4     - POST /api/process ? 202 + job_id (fast-fail 4xx still happens at submit time)
5     - GET /api/jobs/{id} ? {status, progress, result, reports, ...}
6     - job `status` = EXECUTION state (queued|running|done|failed); the pipeline's own
7       verdict (e.g. failed validation) lives in result["status"].
8     - worker crash restores pre-run segments (destroy-AFTER-confirm) and marks the job
failed
9     - startup sweep (fail_stale_jobs) terminalizes jobs orphaned by a restart
10    - migration: fresh DB lands at user_version 2 with a jobs table; v1 DBs upgrade in place
11    """
12
13    import json
14    import sqlite3
15    import time
16    from concurrent.futures import ThreadPoolExecutor
17
18    import pytest
19
20    pytest.importorskip("httpx") # TestClient needs httpx; skip cleanly where absent
21    from fastapi.testclient import TestClient
22
23    import backend.main as m
24    from backend.config.models import AppConfig
25    from backend.infrastructure.database import Database
26    from backend.pipeline.validators.base import ValidationResult
27

```

```

28
29 class _InlineExecutor:
30     """submit() runs the fn synchronously ? deterministic tests, no threads."""
31
32     def submit(self, fn, *args, **kwargs):
33         fn(*args, **kwargs)
34
35
36 @pytest.fixture
37 def env(tmp_path, monkeypatch):
38     db = Database(str(tmp_path / "t.db"))
39     cfg = AppConfig()
40     monkeypatch.setattr(m, "_get_executor", lambda: _InlineExecutor())
41     app = m.create_app(cfg, db)
42     client = TestClient(app)
43     return client, db, tmp_path, monkeypatch
44
45
46 def _video(tmp_path, name="m.mp4"):
47     p = tmp_path / name
48     p.write_bytes(b"not-a-real-video-but-hashable")
49     return p
50
51
52 def _pass():
53     return [ValidationResult(validator_name="x", passed=True, message="ok", details={})]
54
55
56 def _fail():
57     return [
58         ValidationResult(
59             validator_name="blur", passed=False, message="blurry", details={}
60         )
61     ]
62
63
64 class _FakeSeg:
65     def __init__(self, db, cfg):
66         self.db = db
67
68     def process_video(self, video_id, video_path, *a, **k):
69         sid = self.db.add_play_segment(video_id, 0.0, 5.0)
70         return [
71             {
72                 "id": sid,
73                 "start_time": 0.0,
74                 "end_time": 5.0,
75                 "duration": 5.0,
76                 "metadata": {},
77             }
78         ]
79
80
81 class _RaisingSeg:
82     def __init__(self, db, cfg):
83         self.db = db
84
85     def process_video(self, video_id, video_path, *a, **k):
86         self.db.add_play_segment(video_id, 9.0, 10.0) # half-written row pre-crash
87         raise RuntimeError("boom")
88
89
90 # --- migration -----
91
92

```

```

93     def test_migration_fresh_db_is_v3_with_jobs_table(tmp_path):
94         db = Database(str(tmp_path / "fresh.db"))
95         with db._get_connection() as conn:
96             # v9 = async compile (kind/params); v8 = compute-picker; v7 = UI locale; v6 =
cost ledger;
97             # jobs at v2/v3, user_models v5.
98             assert conn.execute("PRAGMA user_version").fetchone()[0] == 9
99             cols = {r["name"] for r in conn.execute("PRAGMA table_info(jobs)").fetchall()}
100         assert {

```

38. user (2026-07-02T04:52:37.253Z)

```

700         if sc is None:
701             return {}
702         return sc.model_dump() if hasattr(sc, "model_dump") else dict(sc)
703
704
705     def _asset_dest_uri(medium_uri: str, video_id: str, ext: str) -> str:
706         """The durable, MD5-keyed destination inside the chosen medium:
707         ``<medium>/<video_id><ext>``. """
708         return f"{medium_uri.rstrip('/')}/{video_id}{ext}"
709
710     def _medium_local_dir(medium_ref, storage_settings: Dict[str, Any]) -> Optional[str]:
711         """The LOCAL directory a ``put`` into this medium would write into (so a free-space
BLOCK can be
712         enforced, D11), or ``None`` for a cloud bucket (which is cost/quota-warned, never
blocked).
713         ``local`` ? the folder; ``gdrive`` ? the resolved path under the Drive mount. """
714         if medium_ref.backend == "local":
715             return medium_ref.key or "."
716         if medium_ref.backend == "gdrive":
717             from backend.storage.capabilities import _gdrive_mount_root
718
719             root = _gdrive_mount_root(storage_settings)
720             if not root:
721                 return None
722             rel = medium_ref.key.lstrip("/")
723             return os.path.join(root, *rel.split("/")) if rel else root
724         return None # s3 / gcs ? a bucket has no readable local capacity
725
726
727     @app.post("/api/assets", status_code=201)
728     def create_asset(req: AssetRequest, request: Request):
729         """Store a video into the user's chosen medium and register it as a library asset
(P4, D14).
730
731         SELF-HOST posture only (D7): refused with 403 when ``security.allowed_video_root``
is set ? the
732         hosted alpha jails inputs and has no per-user medium credentials, so it stays
upload-to-local.
733         Edge-auth mirrors ``/api/process``. Exactly one source: an ``uploaded_path`` (a
local file from
734         the chunked upload) OR a ``source_uri`` (a video already at rest in a medium). An
uploaded file +
735         a REMOTE medium ? a durable, MD5-keyed copy is ``put`` into the medium while the
local upload
736         stays as the working copy; a ``local`` medium keeps the upload in place (it's
already on this box);
737         a pasted at-rest URI is registered as-is (no copy, D14b). The asset is keyed by the
whole-file
738         MD5 (``compute_video_id``, D3) so it joins the pipeline by content hash, and appears
in
739         ``GET /api/videos``. """
740         try:

```

```

741         owner_id = edge_auth.resolve_tenant(request.headers, request.app.state.config)
742     except edge_auth.EdgeAuthError as e:
743         raise HTTPException(status_code=401, detail=f"edge auth: {e}") from e
744
745     cfg = request.app.state.config
746     sec = getattr(cfg, "security", None)
747     if getattr(sec, "allowed_video_root", None):
748         raise HTTPException(
749             status_code=403,
750             detail=(
751                 "storing into a custom medium is disabled on this hosted deployment "
752                 "(security.allowed_video_root is set); upload to local instead."
753             ),
754         )
755     if bool(req.uploaded_path) == bool(req.source_uri):
756         raise HTTPException(
757             status_code=400,
758             detail="provide exactly one of 'uploaded_path' or 'source_uri'.",
759         )
760     if req.uploaded_path and not req.medium_uri:
761         raise HTTPException(
762             status_code=400, detail="'medium_uri' is required when uploading a file."
763         )
764
765     storage_settings = _storage_settings(cfg)
766
767     # 1) Resolve a LOCAL working copy of the source + its whole-file MD5 (the join key,
768     # D3).
769     if req.uploaded_path:
770         try:
771             validate_input_path(req.uploaded_path, allowed_root=None)
772             src_ref = storage.resolve_input_ref(req.uploaded_path, storage_settings)
773         except ValueError as e:
774             raise HTTPException(status_code=400, detail=str(e)) from e
775         if src_ref.backend != "local" or not os.path.isfile(src_ref.key):

```

39. user (2026-07-02T04:52:38.200Z)

```

1545: def _resolve_compile(
1546-     db, cfg, video_id: str, segment_ids: List[int], output_filename: str
1547- ):
1548-     """Shared compile resolution ? run at SUBMIT (fast-fail 4xx) AND in the WORKER (re-
1549-     resolve, since a
1550-     row could change between submit and run). Verifies the video exists, the requested
1551-     segments match
1552-     and survive the ``stitching.min_shot_count`` floor (segments WITHOUT a known shot_count
1553-     are exempt
1554-     so manual/legacy picks can't silently vanish), and the output name is traversal-safe.
1555-     Raises
1556-     ``_CompileError(404/400)``. Returns ``(video_row, kept_segments, safe_out_name)``.
1557-     """
1558-     video = db.get_video(video_id)
1559-     if not video:
1560-         raise _CompileError(404, "Indexed video record not found.")
1561-     # Reject path traversal / absolute names (os.path.join would otherwise write anywhere
1562-     # on disk).
1563-     try:
1564-         out_name = safe_output_filename(output_filename)
1565-     except ValueError as e:
1566-         raise _CompileError(400, str(e)) from e
1567-
1568-     all_segments = db.query_play_segments(video_id, {})
1569-     matched = [s for s in all_segments if s["id"] in segment_ids]
1570-     if not matched:
1571-         raise _CompileError(400, "No matching play segments found to stitch.")
1572-

```

```

1567-     stitching = getattr(cfg, "stitching", None) or StitchingConfig()
1568-     kept = []
1569-     for s in matched:
1570-         meta = s.get("metadata") or {}
1571-         sc = meta.get("shot_count") if isinstance(meta, dict) else None
1572-         try:
1573-             if sc is not None and int(sc) < stitching.min_shot_count:
1574-                 continue
1575-         except (TypeError, ValueError):
1576-             pass # unparseable count -> treat as unknown, keep
1577-         kept.append(s)
1578-     if len(kept) < len(matched):
1579-         logger.info(
1580-             "[compile] stitching.min_shot_count=%s: dropped %d/%d segments below the
floor.",
1581-             stitching.min_shot_count,
1582-             len(matched) - len(kept),
1583-             len(matched),
1584-         )
1585-     if not kept:
1586-         raise _CompileError(
1587-             400,
1588-             f"All {len(matched)} matching segments fall below "
1589-             f"stitching.min_shot_count={stitching.min_shot_count}.",
1590-         )

```

40. user (2026-07-02T04:52:45.964Z)

```

1     import argparse
2     import csv
3     import json
4     import logging
5     import math
6     import os
7     import shutil
8     import socket
9     import sys
10    import tempfile
11    import threading
12    import uuid
13    import webbrowser
14    from importlib.metadata import PackageNotFoundError
15    from importlib.metadata import version as _pkg_version
16    from urllib.parse import quote
17
18    import uvicorn
19    from dotenv import load_dotenv
20
21    # PACKAGING_PLAN.md P0a: prefer an app-home .env ONLY when RALLY_APP_HOME is set (the
packaged
22    # launcher pins it); otherwise env_file_to_load() returns None and load_dotenv(None)
keeps the
23    # exact historical frame-relative walk-up ? byte-identical to the pre-P0a bare
load_dotenv().
24    # stdlib-only import ? keeps this pre-pydantic line light.
25    from backend.utils.app_home import env_file_to_load, resolve_output_dir
26
27    load_dotenv(env_file_to_load())
28
29    # Centralized logging setup for the indexer (hot paths use logger, CLI output uses print
for direct user-facing messages).
30    logging.basicConfig(
31        level=logging.INFO,
32        format="%(%asctime)s [%(levelname)s] %(name)s: %(message)s",
33        datefmt="%H:%M:%S",

```

```

34 )
35 logger = logging.getLogger(__name__)
36 from typing import Any, Dict, List, Optional, Tuple
37
38 from fastapi import FastAPI, HTTPException, Request
39 from fastapi.middleware.cors import CORSMiddleware
40 from fastapi.responses import FileResponse, HTMLResponse, JSONResponse
41 from fastapi.staticfiles import StaticFiles
42 from pydantic import BaseModel
43 from starlette.middleware.trustedhost import TrustedHostMiddleware
44
45 from backend import (
46     billing, # M8: per-job cost ledger (USD internal / INR display)
47     storage, # M2: URI-aware input acquisition (local = identity passthrough)
48 )
49 from backend.api import (
50     auth as edge_auth, # M9: Cloudflare Access edge-auth (default-OFF)
51 )
52 from backend.config import (
53     AppConfig,
54     StartupCheckError,
55     StitchingConfig,
56     enforce_startup_checks,
57     write_light_default_config, # P2b: batteries-included first-run config seeding
58 )
59 from backend.config import ( # new typed config
60     load_config as load_app_config,
61 )
62 from backend.infrastructure.database import Database
63 from backend.pipeline.segmenters.base import segmenter_registry
64 from backend.pipeline.stitcher.compiler import VideoCompiler
65 from backend.pipeline.stitcher.watermark_i18n import localize_watermark
66 from backend.pipeline.validators.base import registry
67 from backend.providers import provider_registry
68 from backend.reporting import emit_indexer_report
69 from backend.results import Failure # standardized error/result contract
70 from backend.storage.capabilities import (
71     storage_capabilities, # P3: secret-free "which media can this box use" report
72 )
73 from backend.utils.ffmpeg import ( # P2a: single ffmpeg/ffprobe resolver
74     ffmpeg_bin,
75     ffprobe_bin,
76 )
77 from backend.utils.freespace import require_free_bytes # HTTP-507 disk guard (P2, D11)
78 from backend.utils.hashing import compute_video_id # canonical file-identity hashing
79 from backend.utils.paths import (
80     output_base,
81     resolve_under_root,
82     safe_output_filename,
83     validate_input_path,
84 )
85 from backend.utils.redact import (
86     redact_secrets, # P0b: mask secret-shaped fields in /api/config
87 )
88 from backend.utils.single_instance import acquire_lock # P0e: one instance per database
89
90 app = FastAPI(title="Sports Highlight Indexer API")
91
92
93 def create_app(config: AppConfig, db: Database) -> FastAPI:
94     """Factory: inject shared state via ``app.state`` so route handlers access it
95     through
96     ``request.app.state`` ? eliminating the Optional module-level globals that forced
97     ``backend.main`` and ``backend.api.job_worker`` into mypy quarantine.

```



```

98     Also sets the module-level ``global_config`` / ``db_instance`` for transitional
99     compatibility (helpers and tests still reference them). Route handlers use
100     ``request.app.state`` exclusively."""
101     global global_config, db_instance
102     global_config = config
103     db_instance = db
104     app.state.config = config
105     app.state.db = db
106     return app
107
108
109     def _cors_origins_from_env(env: Optional[Dict[str, str]] = None) -> List[str]:
110         """M9: CORS allow-origins from the ``RALLY_CORS_ALLOW_ORIGINS`` env var (comma-
111         separated),
112         default ``["*"]``. Read from the ENV (not a cwd ``config.json``) so importing
113         ``backend.main``
114         does zero filesystem I/O ? a hosted/multi-tenant deploy sets this to lock CORS to
115         its origin."""
116         src = os.environ if env is None else env
117         raw = (src.get("RALLY_CORS_ALLOW_ORIGINS", "") or "").strip()
118         origins = [o.strip() for o in raw.split(",") if o.strip()]
119         return origins or ["*"]
120
121     # CORS for the web UI. allow_origins='' with allow_credentials=True is rejected by
122     # browsers per the
123     # fetch spec + is unsafe to carry into multi-tenant; this tool uses no cookies, so
124     # credentials stay

```

41. user (2026-07-02T04:52:59.047Z)

```

1     [mypy]
2     # Lenient, GREEN baseline (quality sweep AI-4): the gate protects the clean
3     # modules and pins an explicit quarantine list below ? not a typing crusade.
4     # Ratchet: when you touch a quarantined module, try removing its entry.
5     ignore_missing_imports = True
6     # backend/ is a namespace package (no root __init__.py) ? required so mypy
7     # maps files to backend.* module names instead of duplicating top-level names.
8     explicit_package_bases = True
9
10    # --- Quarantine (known-loose modules, each with a reason) ---
11
12    # Optional module globals (db_instance/global_config set by create_app() / main()).
13    # Route handlers use request.app.state exclusively; helpers take explicit params.
14    # backend.api.job_worker lifted from quarantine (now takes explicit config/db params).
15    [mypy-backend.main]
16    ignore_errors = True
17
18    # Worker dispatcher: composes the SAME validator/segmenter/report building blocks as
19    # backend.main, which are typed `config: dict` but receive the AppConfig migration shim
20    # at runtime (the .get() compat layer). Shares backend.main's config-shim quarantine
21    # until the typed-config migration removes the dict?AppConfig duality everywhere.
22    [mypy-backend.worker.__main__]
23    ignore_errors = True
24
25    # yolo_hybrid: the fusion-P1 extraction moved its loose torchvision/cv2 inline code
26    # to detectors/person_detector.py (which IS type-clean, no quarantine). What remains
27    # here is pre-existing orchestration debt unrelated to that move (implicit-Optional
28    # process_video defaults, the Success|Failure ABC-override mismatch, a Future[...
29    # annotation) ? ratchet those out in a focused typing pass, not here.
30    # (wasb_hybrid/tracknet_hybrid left the quarantine with the AI-7 collapse.)
31    [mypy-backend.pipeline.segmenters.yolo_hybrid]
32    ignore_errors = True
33
34    [mypy-backend.pipeline.segmenters.gemini]

```

```
35     ignore_errors = True
36
37     # Registry builds cls(None, {}) probes; typing it honestly needs a Protocol.
38     [mypy-backend.pipeline.segmenters.base]
39     ignore_errors = True
40
```

42. user (2026-07-02T04:52:59.471Z)

```
1     # Lint-only lane (no `ruff format` enforcement ? a whole-repo reformat would
2     # bury blame history for zero behavior value).
3     #
4     # Selected: pyflakes + pycodestyle errors (ruff defaults) + pycodestyle
5     # warnings (W) + bugbear (B) + isort (I ? consistent import ordering).
6     target-version = "py38"
7
8     # scratch/ = throwaway debugging scripts that may hit real APIs; it is also
9     # excluded from pytest collection (see pytest.ini).
10    extend-exclude = ["scratch"]
11
12    [lint]
13    extend-select = ["W", "B", "I"]
14
15    [lint.per-file-ignores]
16    # Entrypoint scripts deliberately run setup (load_dotenv, logging config,
17    # sys.path) before importing backend modules ? import order is load-bearing.
18    "backend/main.py" = ["E402"]
19    "scripts/run_phase3_eval.py" = ["E402"]
20    "scripts/run_process_and_stitch.py" = ["E402"]
21
```

43. assistant (2026-07-02T04:53:06.251Z)

I have the full picture. Now setting up the isolated worktree per the branching safety protocol.

44. user (2026-07-02T04:53:12.403Z)

```
Preparing worktree (new branch 'fix/reclaim-input-scratch')
branch 'fix/reclaim-input-scratch' set up to track 'origin/master'.
HEAD is now at e410136 api: /api/videos surfaces fps + resolution (from the stored ffprobe
details) (#413)
fix/reclaim-input-scratch
e410136 api: /api/videos surfaces fps + resolution (from the stored ffprobe details) (#413)
```

45. assistant (2026-07-02T04:53:22.959Z)

Let me check pytest.ini for exclusion patterns, then start implementing.

46. user (2026-07-02T04:53:24.466Z)

```
1     [pytest]
2     # When `pytest` is invoked without an explicit path, only collect from tests/.
3     # scratch/ holds debugging scripts that hit real APIs (Gemini upload, etc.)
4     # and should never run as part of the regular test suite.
5     testpaths = tests
6
```

47. user (2026-07-02T04:53:37.551Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

48. user (2026-07-02T04:53:45.151Z)

ok

49. user (2026-07-02T04:53:53.603Z)

```
1      """Pluggable storage layer: local (default) + S3/GCS first-class buckets + Google Drive.
2
3      URI-driven (see ``ref.parse_uri``). The package import is cheap ? cloud SDKs are
imported
4      lazily, only when a non-local bucket URI is actually used (install the ``storage-s3`` /
5      ``storage-gcs`` extras). The default ``local`` backend needs no third-party deps and
preserves
6      today's behaviour byte-for-byte; ``gdrive://`` likewise needs none ? it reads a locally
MOUNTED
7      Google Drive (Google Drive for Desktop) as plain files.
8
9      Worker/dispatcher entry points speak URIs via the thin helpers ``fetch`` / ``put`` /
10     ``list_uris`` / ``get_storage``; ``local://`` (and bare paths) make ``fetch`` an
identity
11     passthrough. See ``docs/DECOUPLED_COMPUTE``.
12     """
13
14     from __future__ import annotations
15
16     import os
17     import tempfile
18     from typing import Any, Dict, Iterator, Optional
19
20     from backend.storage.base import StorageBackend
21     from backend.storage.ref import StorageRef, StorageStat, parse_uri, resolve_input_uri
22
23     __all__ = [
24         "StorageBackend",
25         "StorageRef",
26         "StorageStat",
27         "parse_uri",
28         "resolve_input_uri",
29         "get_storage",
30         "fetch",
31         "put",
32         "list_uris",
33         "acquire_local_input",
34         "input_is_local",
35         "resolve_input_ref",
36     ]
37
38
39     def _storage_dict(storage_config: Any) -> Dict[str, Any]:
40         """Coerce a StorageConfig (Pydantic) / dict / None into a plain settings dict."""
41         if storage_config is None:
42             return {}
43         if hasattr(storage_config, "model_dump"):
44             return storage_config.model_dump() # type: ignore[no-any-return]
45         return dict(storage_config)
46
47
48     def resolve_input_ref(raw: str, storage_config: Any = None) -> StorageRef:
49         """Resolve a user-supplied input (a path or storage URI) into a ``StorageRef`` under
the input
50         config (``input_backend``/``input_root``). Raises ``ValueError`` for an unsupported
URI scheme ?
51         callers at the API/CLI boundary catch that and return a clean client error (never a
500).
52
53         Use ``ref.backend`` to branch local-vs-remote and ``ref.key`` for the resolved LOCAL
path (e.g.
54         ``local://?`` is stripped to the bare path), rather than stat-ing the raw URI
string."""
```

```

55     sc = _storage_dict(storage_config)
56     uri = resolve_input_uri(
57         raw,
58         input_backend=sc.get("input_backend", "local"),
59         input_root=sc.get("input_root", ""),
60     )
61     return parse_uri(uri)
62
63
64     def input_is_local(raw: str, storage_config: Any = None) -> bool:
65         """True if ``raw`` resolves (under the input config) to a LOCAL filesystem path ?
i.e. a
66         plain on-disk file the caller can stat/validate directly. False for a bucket / Drive
URI
67         that must be fetched first. Raises ``ValueError`` for an unsupported scheme (same as
68         ``resolve_input_ref``)."""
69         return resolve_input_ref(raw, storage_config).backend == "local"
70
71
72     def acquire_local_input(
73         raw: str, storage_config: Any = None, *, scratch_parent: "str | None" = None
74     ) -> str:
75         """Resolve ``raw`` (a path or storage URI) to a LOCAL file path ready for
processing.
76
77         Local / bare paths are returned UNCHANGED (identity passthrough, no copy ? today's
behaviour
78         byte-for-byte). A bucket / Drive source (``gs://``/``s3://``/``gdrive://``, or a
bare name under a
79         non-local ``input_root``/``input_backend``) is downloaded into a fresh scratch dir
and the local
80         path returned. The scratch copy is left under the system temp dir (a pure
intermediate,
81         regenerable from the source) and relies on OS temp reclamation ? the same pull
pattern the
82         worker's ``cmd_infer`` already ships."""
83         sc = _storage_dict(storage_config)
84         ref = resolve_input_ref(raw, sc)
85         if ref.backend == "local":
86             return (
87                 ref.key
88                 ) # identity ? no temp, no copy, byte-identical to passing the path directly
89         scratch = tempfile.mkdtemp(prefix="rally-input-", dir=scratch_parent)
90         return fetch(ref.uri, os.path.join(scratch, ref.name or "in.mp4"), config=sc)
91
92
93     def _backend_class(name: str):
94         """Lazy-resolve a backend class so cloud SDKs aren't imported until needed."""
95         if name == "local":

```

50. user (2026-07-02T04:54:07.672Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/storage/__init__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

51. user (2026-07-02T04:54:27.876Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/storage/__init__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

52. assistant (2026-07-02T04:54:35.589Z)

Now the sweep module.

53. assistant (2026-07-02T04:54:58.545Z)

""Age-based startup sweep for orphaned fetched-input scratch dirs (R1-17).

``acquire\_local\_input`` downloads every remote input (``gs:///``s3:///``gdrive:///``) into a fresh ``rally-input-\*`` mkdtemp, and the worker's ``cmd\_infer`` uses ``rally-worker-\*`` the same way. The per-job scope (``acquire\_local\_input\_scoped``) reclaims those on the normal success/failure paths, but a crash / SIGKILL / pre-scope leak leaves multi-GB video copies under the system temp dir FOREVER on the shipping platform (Windows never sweeps ``%TEMP%`` ? the old "relies on OS temp reclamation" premise was false). This sweep runs once in the server/CLI boot path and deletes only:

- **directories** (never plain files),
- whose basename carries one of OUR ownership-marker prefixes (``SWEEP\_PREFIXES``),
- directly under the system temp dir (or an explicit ``temp\_root``),
- older than the age threshold (default 3 days ? far beyond any legitimate in-flight job, so a running fetch/segment can never be swept out from under itself).

DEFAULT-ON with a config kill-switch (``storage.scratch\_sweep\_enabled=false``) and an age knob (``storage.scratch\_sweep\_max\_age\_days``). Every removal is logged. Never raises ? a janitor must not be able to break boot.

""

```
from __future__ import annotations
```

```
import logging
```

```
import os
```

```
import shutil
```

```
import tempfile
```

```
import time
```

```
from typing import Any, List, Optional
```

```
from backend.storage import SCRATCH_INPUT_PREFIX
```

```
logger = logging.getLogger(__name__)
```

```
#: Ownership markers: only dirs whose basename starts with one of these are ever touched.
```

```
#: ``rally-input-`` = acquire_local_input fetch scratch; ``rally-worker-`` = the worker
```

```
#: ``cmd_infer`` scratch (``backend/worker/__main__.py``).
```

```
SWEEP_PREFIXES = (SCRATCH_INPUT_PREFIX, "rally-worker-")
```

```
#: Default age threshold (days) before an orphaned scratch dir is reclaimed.
```

```
DEFAULT_MAX_AGE_DAYS = 3.0
```

```
def sweep_orphan_scratch(
```

```
    temp_root: Optional[str] = None,
```

```
    *,
```

```
    max_age_days: float = DEFAULT_MAX_AGE_DAYS,
```

```
    now: Optional[float] = None,
```

```
) -> List[str]:
```

```
    """Delete orphaned scratch dirs under ``temp_root`` (default: the system temp dir) older
    than ``max_age_days``; return the paths removed. Age = the dir's own mtime (set at mkdtemp
    and bumped when a fetch lands a file in it), so anything younger than the threshold ? in
    particular every in-flight job ? is left alone. Best-effort per entry and never raises:
    an unstatable/undeletable entry is skipped (and retried at the next boot)."""
```

```
    root = temp_root or tempfile.gettempdir()
```

```
    cutoff = (now if now is not None else time.time()) - max_age_days * 86400.0
```

```
    removed: List[str] = []
```

```
    try:
```

```
        entries = os.listdir(root)
```

```

except OSError as e:
    logger.warning("[SCRATCH-SWEEP] cannot list %s (%s); skipping.", root, e)
    return removed
for name in entries:
    if not name.startswith(SWEEP_PREFIXES):
        continue
    path = os.path.join(root, name)
    try:
        st = os.lstat(path)
    except OSError:
        continue # vanished / unreadable ? not ours to force
    if not os.path.isdir(path) or os.path.islink(path):
        continue # only real directories carry our mkdtemp marker
    if st.st_mtime >= cutoff:
        continue # too young ? could belong to an in-flight job
    try:
        shutil.rmtree(path)
    except OSError as e:
        logger.warning("[SCRATCH-SWEEP] could not remove %s (%s); skipped.", path, e)
        continue
    removed.append(path)
    logger.info(
        "[SCRATCH-SWEEP] removed orphaned scratch %s (age %.1f days).",
        path,
        ((now if now is not None else time.time()) - st.st_mtime) / 86400.0,
    )
return removed

```

```

def maybe_sweep_scratch(config: Any) -> List[str]:
    """Boot-path wrapper: run the sweep under the config knobs. DEFAULT-ON (the sweep only ever
    touches this tool's own prefix-marked dirs, aged well past any live job); kill-switch =
    ``storage.scratch_sweep_enabled=False``; threshold = ``storage.scratch_sweep_max_age_days``.
    Returns the removed paths. Never raises ? a failed sweep is logged and swallowed."""
    try:
        sc = getattr(config, "storage", None)
        if not getattr(sc, "scratch_sweep_enabled", True):
            return []
        max_age = float(
            getattr(sc, "scratch_sweep_max_age_days", DEFAULT_MAX_AGE_DAYS)
        )
        removed = sweep_orphan_scratch(max_age_days=max_age)
        if removed:
            logger.info(
                "[STARTUP] scratch sweep reclaimed %d orphaned dir(s) older than %s day(s).",
                len(removed),
                max_age,
            )
        return removed
    except Exception as e: # noqa: BLE001 ? janitor must never break boot
        logger.warning("[STARTUP] scratch sweep skipped (non-fatal): %s", e)
        return []

```

54. user (2026-07-02T04:54:58.578Z)

File created successfully at: C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/storage/scratch_sweep.py (file state is current in your context ? no need to Read it back)

55. assistant (2026-07-02T04:55:03.171Z)

Now the config knobs.

56. user (2026-07-02T04:55:04.888Z)

```
380     input_root: str = ""
381     # --- Per-video workspace (docs/PER_VIDEO_WORKSPACE.md) ---
382     # One folder per video keyed by the MD5 video_id:
383     <root>/[<workspace_prefix>]/<video_id>/...
384     # so a local deploy and a cloud (bucket) deploy share ONE relative layout (only the
385     root
386     # differs). DEFAULT-OFF: convergence is phased + flipped via a Launch Report (D5).
387     When OFF
388     # the legacy flat layout is used unchanged.
389     single_folder_per_video: bool = False
390     # Durable workspace root URI. "" => precedence: output_root, then output.output_dir.
391     workspace_root: str = ""
392     # Cloud namespace prefix under the root so per-video folders sit tidily beside
393     # videos/ labels/ weights/ (e.g. gs://bucket/workspaces/<video_id>/). "" = root
394     directly.
395     workspace_prefix: str = "workspaces"
396     local: dict[str, Any] = Field(default_factory=dict)
397     s3: dict[str, Any] = Field(default_factory=dict)
398     gcs: dict[str, Any] = Field(default_factory=dict)
399     gdrive: dict[str, Any] = Field(default_factory=dict)
400     # gdrive-api settings (e.g. {"root_folder_id": "<id>"}); read via the
401     hyphen?underscore lookup in
402     # backend.storage.get_storage. NO secrets ? creds come from ADC over the Drive
403     scope.
404     gdrive_api: dict[str, Any] = Field(default_factory=dict)
405
```

57. user (2026-07-02T04:55:16.585Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/config/models.py has been updated successfully. (file state is current in your context ? no need to Read it back)

58. assistant (2026-07-02T04:55:21.845Z)

Now `cloud\_run.\_read\_json`.

59. user (2026-07-02T04:55:29.346Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

60. user (2026-07-02T04:55:38.252Z)

```
200         execution=str(execution or ""),
201     )
202
203     # --- storage-backed completion signal (bucket-poll) ---
204     def _read_marker(self, done_uri: str) -> Optional[dict]:
205         """The done-marker dict if the worker has written it, else None ? poll_until
206         keeps waiting."""
207         ref = parse_uri(done_uri)
208         if not get_storage(ref.backend, config=self._storage_config).exists(ref):
209             return None
210         return (
211             self._read_json(done_uri) or {}
212         ) # empty dict still terminates the poll (present == done)
213
214     def _read_json(self, uri: str) -> dict:
215         # Download to a real local dst FIRST: the cloud backends (gcs/s3) REQUIRE a dst
216         # ? a None dst
217         # raises ValueError ? unlike local's identity passthrough. Without an explicit
```

```

dst this read
216         # would silently return {} for every gs://|s3:// object (the only backends this
shim targets).
217         try:
218             ref = parse_uri(uri)
219             dst = os.path.join(
220                 tempfile.mkdtemp(prefix="cr-read-"), ref.name or "obj.json"
221             )
222             local = fetch(uri, dst, config=self._storage_config)
223             with open(local, "r", encoding="utf-8") as f:
224                 data: Any = json.load(f)
225             return data if isinstance(data, dict) else {}
226         except Exception:
227             LOG.warning("cloud-run: could not read %s", uri, exc_info=True)
228             return {}
229

```

61. user (2026-07-02T04:55:51.608Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/compute/cloud_run.py has been updated successfully. (file state is current in your context ? no need to Read it back)

62. user (2026-07-02T04:56:03.512Z)

```
21:requires-python = ">=3.10"
```

63. user (2026-07-02T04:56:11.936Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/compute/cloud_run.py has been updated successfully. (file state is current in your context ? no need to Read it back)

64. assistant (2026-07-02T04:56:54.134Z)

Let me check the ReportConfig default ? it matters for whether reports land in scratch.

65. user (2026-07-02T04:56:56.570Z)

```

272-     default_highlights_name: str = "highlights.mp4"
273-     db_name: str = "sports_indexer.db"
274-     # Directory where the DB, highlight reels, and processed clips are written.
275-     # The CLI's --output-dir overrides this at startup (see backend/main.py::main).
276-     output_dir: str = "output"
277-
278-
279-class WatermarkConfig(BaseModel):
280-     """Branding overlay burned into each highlight clip during the per-clip re-encode
--
365-     **secrets are never stored here**, only endpoints / regions / file paths / source-
selectors
366-     (boto3/google auth chains supply credentials). See ``backend/storage/``.
367-
368-     backend: Literal["local", "s3", "gcs", "gdrive", "gdrive-api"] = "local"
369-     # Output root for produced artifacts. "" => fall back to output.output_dir (non-
breaking).
370-     output_root: str = ""
371-     # --- Input side (COMPUTE_DECOUPLED_SERVING M2; parallel to the output
`backend`/`output_root`).
372-     # Default ``local`` / "" keeps today's behaviour: a bare path / ``local://`` URI is read
IN PLACE
373-     # (identity, no copy). Set these to read inputs from a bucket / mounted Drive ? a

```



```

bare/relative
--
391- # it only ever deletes this tool's own prefix-marked dirs, aged well past any live job
392- # (Windows never sweeps %TEMP% on its own). Kill-switch: set
scratch_sweep_enabled=false.
393- scratch_sweep_enabled: bool = True
394- scratch_sweep_max_age_days: float = 3.0
395: # Durable workspace root URI. "" => precedence: output_root, then output.output_dir.
396- workspace_root: str = ""
397- # Cloud namespace prefix under the root so per-video folders sit tidily beside
398- # videos/ labels/ weights/ (e.g. gs://bucket/workspaces/<video_id>/). "" = root
directly.
399- workspace_prefix: str = "workspaces"

```

66. user (2026-07-02T04:57:05.352Z)

```

384: # differs). DEFAULT-OFF: convergence is phased + flipped via a Launch Report (D5). When
OFF
482: # the alias never deprecates underneath us. Eval reports record the run date alongside.

```

67. user (2026-07-02T04:57:17.197Z)

```

1161         notify = progress or (lambda stage: None)
1162
1163         notify("hashing")
1164         # M2: resolve the input to a LOCAL path. A bare/local path is returned unchanged
(identity ?
1165         # byte-identical to today); a bucket/Drive URI is fetched to scratch first. The
original
1166         # req.video_path (the source URI) is kept for the DB record + report; only the file-
reading
1167         # consumers (hash / validators / segmenter) use the resolved local path.
1168         local_video = storage.acquire_local_input(
1169             req.video_path, getattr(config, "storage", None)
1170         )
1171         video_id = compute_video_id(local_video)
1172         was_reingest = db.get_video(video_id) is not None
1173
1174         # 1. Run pluggable validators (with-context variant surfaces ffprobe_meta for the
report)
1175         notify("validating")
1176         logger.info(f"Running validations for {req.video_path}...")
1177         val_results, val_context = registry.run_all_with_context(local_video, config)
1178         failures = [r for r in val_results if not r.passed]
1179
1180         # typed AppConfig: attribute access for the default segmenter (with safe fallback)
1181         default_seg = getattr(
1182             getattr(config, "indexing", None), "default_segmenter", "motion"
1183         )
1184         seg_name = req.segmenter or default_seg
1185         segmenter_actual = None
1186         segmentation_ran = False
1187         # Compute-path PROVENANCE (the validation method): which backend actually ran the
DETECT ?
1188         # "in_process" once the local segmenter is reached, set on the response as
"cloud_run" by
1189         # _maybe_dispatch_remote when it dispatches. None ? nothing ran (validation/config
fail).
1190         compute_actual: Optional[str] = None
1191         response: Optional[Dict[str, Any]] = None
1192         try:
1193             if failures:
1194                 # add_video is an UPSERT (no cascade) ? a failed re-validation only updates
1195                 # status; it no longer destroys the video's prior good segments.

```

68. user (2026-07-02T04:57:35.601Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

69. assistant (2026-07-02T04:57:42.260Z)

Now the end of `\_execute\_processing`'s finally block.

70. user (2026-07-02T04:57:47.301Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

71. assistant (2026-07-02T04:57:55.841Z)

Site 2: `\_execute\_compile`.

72. user (2026-07-02T04:58:04.062Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

73. assistant (2026-07-02T04:58:08.947Z)

Site 3: `create\_asset`.

74. user (2026-07-02T04:58:13.356Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

75. user (2026-07-02T04:58:27.448Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-scratch-reclaim/backend/main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

76. assistant (2026-07-02T04:58:28.177Z)

You've hit your session limit · resets 1:10pm (Asia/Calcutta)