

claude-windows / c1fcb4aa-537d-41fd-858e-53f93349bf5a

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\subagents\workflows\wf_41c8dc8a-8a4\agent-ab50ae8590de2ddef.jsonl`
- SHA-256: `c7308b0f69463d451270a9e9c6c0176fbb1be0405affc1290ba117c871aa453e`
- Source modified: `2026-07-02T03:04:11+00:00`
- Imported at: `2026-07-05T16:48:24+00:00`
- project: `wf_41c8dc8a-8a4`
- session_id: `c1fcb4aa-537d-41fd-858e-53f93349bf5a`
```

Transcript

1. user (2026-07-02T03:02:49.419Z)

You are an adversarial verifier in a tech-debt audit. Your default stance: REFUTE. Today is 2026-07-02.
A finder claims the following about the repo at C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector:

```
{
  "title": "Uncommitted .gitignore hunk (artifacts/) stale ~3 weeks and self-contradictory vs the ADR it cites",
  "category": "hygiene",
  "severity": "P2",
  "file": ".gitignore",
  "line": 59,
  "evidence": "git diff shows an uncommitted hunk appending lines 56-59: '# Model artifact bundles: weights are NEVER committed (private Releases); the reviewable surface is each bundle's manifest.json (ADR-008).' + 'artifacts/'. File mtime is Jun 10 15:53 ? same day .gitignore was last committed (e4dd8f1, 2026-06-10) ? so this WIP has sat uncommitted ~3 weeks across sessions. ADR-008 does not exist in this repo (docs/DECISIONS.md has only ADR-001/ADR-002); it lives in the indexer at badminton-highlight-indexer/docs/archives/decisions/DECISIONS.md:388 ('Owned-model topology'). Internal contradiction: a blanket 'artifacts/' ignore also hides the manifest.json the comment declares reviewable (same contradiction exists in the indexer: its .gitignore:53 ignores artifacts/ wholesale while its CODE_MAP.md:35 claims 'manifest committed'). No artifacts/ dir exists in this repo yet.",
  "proposed_fix": "Either commit the hunk as a tiny PR with a pattern that actually matches the stated intent (e.g. 'artifacts/**' + '!artifacts/', '!artifacts/**/', '!artifacts/**/manifest.json'), or discard it if model artifact bundles will only ever live in the indexer repo (per ADR-008 / indexer CODE_MAP, artifacts/<name>/<ver>/ is an indexer-repo location). Owner should confirm which repo hosts bundles before committing.",
  "effort": "S",
  "tracked_in": ""
}
```

Repo root candidates: C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer, C:/Users/avidu/Projects/khelsutra-guru/khelsutra, C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector, C:/Users/avidu/Projects/khelsutra-guru/rally-annotator, C:/Users/avidu/Projects/khelsutra-guru/sports-obsreport (pick from the finding's repo field; non-git dirs: 'C:/Users/avidu/Projects/khelsutra-guru/Annotation Setup', C:/Users/avidu/Projects/khelsutra-guru/khelsutra-screencast).

VERIFY against the CURRENT working tree (just pulled to origin master/main): (1) does the cited code actually exist and behave as claimed at file:line? (2) was it already fixed by a recent commit (git log --oneline -15 -- <file>)? (3) is the severity honest - construct the concrete failure scenario or downgrade; (4) is the proposed fix sane and correctly scoped; (5) is it owner-gated per: alpha/serving go-live, Studio P10/P11 corpus-promotion/training, product/strategy/licensing, real cloud spend, counsel ToS/DPA?

HARD RULES: READ-ONLY (no file modifications, no installs, no test-suite runs). gh read-only

commands allowed. Your reasoning MUST cite file:line you actually read. CONFIRMED requires you personally saw the defect/debt in current code; if you cannot reproduce the claim from the code, say REFUTED; if a commit fixed it, ALREADY_FIXED.

2. user (2026-07-02T03:02:57.203Z)

```
M .gitignore
?? sports-data-collector/
---DIFF---
warning: in the working copy of '.gitignore', LF will be replaced by CRLF the next time Git
touches it
diff --git a/.gitignore b/.gitignore
index 52dac92..ac1420f 100644
--- a/.gitignore
+++ b/.gitignore
@@ -53,3 +53,7 @@ data/temp/

# Local DB backups (e.g. data/sports_metadata.backup-YYYY-MM-DD.db) ? never commit
data/*.backup-*.db
+
+# Model artifact bundles: weights are NEVER committed (private Releases); the
+# reviewable surface is each bundle's manifest.json (ADR-008).
+artifacts/
---LOG---
e4dd8f1 chore(corpus): reclassify 23 ShuttleSet broadcast videos as non-commercial (owner-
approved) (#17)
038424d feat(reporting): per-video observability report (collector)
d4327c5 chore: address code-review tech debt (config, parsing, deps, fixes)
854b077 Initial commit: sports data collector & curator
```

3. user (2026-07-02T03:02:58.463Z)

```
adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)
6f67d65 feat(import): callable import_annotations + a maturity gate (P9) (#52)
c281f82 feat(export): add source file MD5 to the eval/training manifest (P9) (#51)
master
-rw-r--r-- 1 avidu 197609 1292 2026-06-10 15:53 .gitignore
```

4. assistant (2026-07-02T03:03:04.663Z)

The diff exists as claimed. Now let me verify the ADR-008 claims and dates.

5. user (2026-07-02T03:03:07.864Z)

```
e4dd8f1 2026-06-10 14:42:02 +0530 chore(corpus): reclassify 23 ShuttleSet broadcast videos as
non-commercial (owner-approved) (#17)
---ADR search collector---
docs/annotation/ROORA_RUNBOOK.md:8:ADR-002 in [`.../DECISIONS.md`](../DECISIONS.md).
docs/annotation/ROORA_RUNBOOK.md:197:grid. Full decision record: ADR-002 in
[`.../DECISIONS.md`](../DECISIONS.md);
docs/archives/CODE_HEALTH_PLAN.md:6:> **Relation to existing docs:** peer-of `docs/DECISIONS.md`
(ADR-001, ADR-002); the de-dup work items (T1/A4) are explicitly **gated behind an ADR-001
revisit** (see D2).
docs/archives/CODE_HEALTH_PLAN.md:8:> **Honesty note:** every P0/P1 claim below is tagged
`[verified]` (checked against current `master` source) or `[researched]` (needs sign-off).
ADR-001 is `[verified]`; the refactor items that conflict with it are marked ? gated.
docs/archives/CODE_HEALTH_PLAN.md:14:A full code-health audit of `sports-data-collector`
surfaced **7 correctness/license-integrity bugs (P0)**, **10 tech-debt / non-intrusive
improvements (P1)**, and **6 structural/architectural flags (P2)**. The plan sequences work
**safest-first** (tiny, isolated, fully-tested wins ? larger refactors) so we warm up before
touching the duplicated per-sport code. The single most important caveat: **items T1 & A4 (the
de-dup refactor) conflict with [ADR-001](DECISIONS.md#adr-001--keep-explicit-per-sport-code-
instead-of-a-sport-registry-for-now), which deliberately keeps per-sport code ? they are owner-
gated behind an ADR revisit, not planned for execution.** The highest-leverage standalone win is
```

B1 (``is_commercial`` not bound to ``license_type``): it is a license-compliance landmine for a product whose commercial-clean licensing is non-negotiable.

docs/archives/CODE_HEALTH_PLAN.md:22:**Read first:** ``README.md``, ``docs/DECISIONS.md`` (ADR-001 + ADR-002), ``docs/PIPELINE_ROLE.md``, ``src/core/db.py``, ``src/core/schemas/base.py``, ``src/parsers/cvat/rally_cvat.py``.

docs/archives/CODE_HEALTH_PLAN.md:29: | D2 | **Honor ADR-001.** Do NOT plan the per-sport de-dup refactor (T1, A4) for execution. List it as owner-gated; only revisit if an ADR-001 trigger fires. | ADR-001 ``[verified]`` | T1/A4 are tracked but ?; the "Future design sketch" in ADR-001 is the starting point if revisited. |

docs/archives/CODE_HEALTH_PLAN.md:35: | D8 | **Tier D triage (owner, 2026-06-24).** T1/A4 (sport registry) **dropped** ? ADR-001 stands. A2/A5/A6 **promoted to active** (independently valuable, no backend needed). A1 (crawler retry) + A3 (read-only DB) **parked** in ``docs/NEXT_STEPS.md`` with explicit criticality triggers ? A1 because the batch-vs-single merge is messy/hard to test thoroughly without the live golden corpus; A3 because it's speculative until a backend exists. | Owner, 2026-06-24 | Resolves Q06/Q07. Tier D is no longer "flag-only" ? A2/A5/A6 are now executable work items below; A1/A3 are deferred with clear unpark conditions. |

docs/archives/CODE_HEALTH_PLAN.md:39:- **[verified]** ADR-001 exists and explicitly keeps per-sport code; its "revisit triggers" are: sport count > ~8, a 4th+ rally-identical sport, or a cross-cutting feature requiring edits to every per-sport branch. **None have fired.** ? T1/A4 stay gated.

docs/archives/CODE_HEALTH_PLAN.md:48: No new architecture is introduced by the P0/P1 work ? each item is a localised fix. The only structural proposal is **A4 (sport registry)**, which is the same idea ADR-001 already sketches and defers; it is **not planned here**.

docs/archives/CODE_HEALTH_PLAN.md:68:- T1/A4 (de-dup) are **not executed** ? the duplication remains until ADR-001 is revisited.

docs/archives/CODE_HEALTH_PLAN.md:132: | T1 / A4 | Per-sport de-dup refactor / sport registry. | **Conflicts with ADR-001** (deliberately keeps explicit per-sport code). T2 already removed the worst schema-level duplication. Revisit only if an ADR-001 trigger fires. |

docs/archives/CODE_HEALTH_PLAN.md:151:- [x] ADR-001 is left untouched (T1/A4 dropped per D8) ? never silently contradicted.

docs/archives/CODE_HEALTH_PLAN.md:167:- ``docs/DECISIONS.md`` ? ADR-001 (per-sport code, ? to T1/A4), ADR-002 (proxy contract).

docs/archives/CODE_HEALTH_PLAN.md:175:- ``2026-06-24`` ? Created from the 2026-06-24 full code-health audit. Docs-only PR (D6); no code items closed. Tier S?D sequencing + ADR-001 gating recorded as decisions D1?D6.

docs/archives/CODE_HEALTH_PLAN.md:180:- ``2026-06-24`` ? **Tier D triaged (D8).** T1/A4 dropped (ADR-001 stands); A2/A5/A6 promoted to active; A1/A3 parked in ``docs/NEXT_STEPS.md``. All Qs resolved. Final 3 items (A2/A5/A6) remain before DONE.

docs/archives/CODE_HEALTH_PLAN.md:181:- ``2026-06-24`` ? **PLAN DONE.** A6 merged (#46); A2 merged (#47); A5 merged (#48). All active Tier D items complete. A1/A3 remain parked in NEXT_STEPS with triggers; T1/A4 dropped per ADR-001. Final: 280 tests / 82% coverage (was 193 / 81%). Status ? DONE; archived to ``docs/archives/``.

docs/DECISIONS.md:8:## ADR-001 ? Keep explicit per-sport code instead of a sport registry (for now)

docs/DECISIONS.md:41:- If the cost rises, ADR-001 should be revisited and superseded.

docs/DECISIONS.md:65:## Checklist: adding a new sport today (under ADR-001)

docs/DECISIONS.md:76:## ADR-002 ? Annotation proxies are generated in-repo (``curate prep-annotation-proxy``), not by a one-off tool

docs/INGESTION_PIPELINE.md:35:## (ADR-002 ? downscaled + scrub-friendly so the vendor can work at frame step 1;

docs/INGESTION_PIPELINE.md:167: the original. Full rationale: ADR-002 in ``docs/DECISIONS.md``.

docs/PIPELINE_ROLE.md:13:- ``badminton-highlight-indexer/docs/DECISIONS.md`` ? ADR-002 (this bridge), ADR-003

docs/PIPELINE_ROLE.md:14: (acquisition policy), ADR-004 (offline learned segmenter).

docs/PIPELINE_ROLE.md:33: **segmentation**, not the frame-level shuttle detector (see the indexer's ADR-002).

docs/PIPELINE_ROLE.md:38: hard-won pitfalls (SABR?144p, Chrome DPAPI, rate-limiting) are in the indexer's ADR-003.

---artifacts dir?---

no artifacts dir

6. user (2026-07-02T03:03:09.006Z)

badminton-highlight-indexer\training__init__.py:1: ""Owned-model TRAINING tree (ADR-008). backend/ must never import this package ?

```

badminton-highlight-indexer\training\requirements-train.txt:1:# Training-only dependencies
(ADR-008): NEVER merge these into the root requirements.txt.
badminton-highlight-indexer\training\requirements-train.txt:12:# Gen-2 (when greenlit; ADR-008
split trigger):
badminton-highlight-indexer\backend\features\__init__.py:1:"""Single-source-of-truth feature
extraction shared by training and serving (ADR-008).
badminton-highlight-indexer\training\README.md:1:# training/ ? the owned-model training tree
(ADR-008)
badminton-highlight-indexer\training\README.md:6:the split happens when the ADR-008 triggers
fire (Gen-2 / ms-swift / cloud training, or a
badminton-highlight-indexer\backend\features\rally_seq.py:3:THE single source of truth for the
Gen-0 TAL feature extraction (ADR-008): the training
badminton-highlight-indexer\training\gen0\harness.py:6:Productionizes the rally_seq prototype
per ADR-008:
badminton-highlight-indexer\backend\eval\calibrate_wasb.py:45:    start,end ? the historical
fork is closed; see gt_loader.py / ADR-008). Keeps this
badminton-highlight-indexer\backend\eval\annotations.py:47:    Delegates to the canonical loader
(gt_loader.py, ADR-008 prereq), which accepts BOTH
badminton-highlight-indexer\backend\eval\gt_loader.py:1:"""Canonical ground-truth interval
loader ? THE one GT/golden-CSV reader (ADR-008 prereq).
badminton-highlight-indexer\backend\eval\rally_seq_proto.py:43:# Feature extraction is single-
sourced in backend/features (ADR-008): training and
badminton-highlight-indexer\docs\NEXT_STEPS.md:51:`docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` (+
ADR-008 for topology). **Latest blow-by-blow:** memory
badminton-highlight-indexer\docs\NEXT_STEPS.md:70:ADR-008 ratified ("repo split driven by
productionization, not isolation"): training lives in-repo (`training/`)
badminton-highlight-indexer\docs\NEXT_STEPS.md:154:- Roadmap:
`docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` · Topology: ADR-008
(`docs/archives/decisions/DECISIONS.md`) ·
badminton-highlight-indexer\docs\COMPUTE_DECOUPLED_SERVING\README.md:52:| **ADR-008** | Owned-
model topology; **train==serve featurizer single-sourcing** | The "core unaffected across
profiles" is the same parity discipline, extended from train/serve to **local/cloud**. |
badminton-highlight-indexer\docs\CODE_MAP.md:31:| Training code (model fitting) | `training/` |
Behind the CI-enforced import wall (ADR-008) ? `backend.main` must not import it. Own
`requirements-train.txt`. |
badminton-highlight-indexer\docs\CODE_MAP.md:35:| A model artifact (weights + manifest) |
`artifacts/<name>/<ver>/` | ADR-008: **manifest committed** (provenance), **weights
gitignored**. |
badminton-highlight-indexer\docs\CODE_MAP.md:177:[Omitted long matching line]
badminton-highlight-indexer\docs\CODE_MAP.md:178:[Omitted long matching line]

```

[Showing results with pagination = limit: 20]

7. assistant (2026-07-02T03:03:23.925Z)

Let me verify the indexer cross-references and check whether the indexer actually tracks any manifest.json under artifacts/.

8. user (2026-07-02T03:03:25.294Z)

```

output_dir/
backend/static/highlights_*.mp4
backend/static/*_compiled.mp4

```

Local generated artifacts (review reels, eval CSVs, e2e/proof outputs) -- regenerable, n
artifacts/

OS generated files

```

.DS_Store
Thumbs.db
desktop.ini
---ADR-008 header---
388:## ADR-008: Owned-model topology ? artifact boundary now, repo split at Gen-2

```

498: failure mode is documented in ADR-008's rationale) with pointer READMEs in `khealsutra`.
691: `ADR-005` (permissive licensing ? engine + image stay MIT/Apache/BSD) . `ADR-008`
(train==serve featurizer single-sourcing ? the parity discipline this extends to local?cloud) .
`ADR-009` (**local-first delivery** ? origin of the truly-local profile) . `ADR-010` (DECOUPLED
DO?GCP) . `ADR-011` (**deferred M3 serving + M4 billing**; overrode "rent nothing / not a
service"; \$06 gates = DPA/consent, collector `md5`, WASB-C3) . `ADR-012` (**Cloud Run+GPU =
serving model**; NVDEC lever) . `ADR-013` (**GCP sole compute stack**) . **? ADR-014** (this).
Non-ADR: realizes `PLATFORM\_ARCHITECTURE.md`'s `ComputeBackend`/`StorageBackend` registries; the
truly-local profile = `VIDEO\_LOCALITY\_MODEL.md` L0 + `DESKTOP\_APP\_PLAN.md` `LocalDesktop`;
informed by the archived `DEPLOY\_REPORT\_GCP\_2026-06-20` (cross-GPU sub-pixel parity finding).
---tracked artifacts---
---artifacts dir---
appliance_design.json
branch-prune-backup-24178a7.txt
cloudrun_e2e_validation
cloudrun_lightup
landing
long_rally_lens_n15.json
long_rally_lens_n15.txt
multishuttle_142
pr245_code_review.md
pr245_fixes_plan.md
pr245_review.json
pr245_review.md
promo
rally_tal_gen0
review
site_deploy
so.json
so.py
tridet.pdf
tridet.txt

9. assistant (2026-07-02T03:03:56.017Z)

All claims verified against the current working tree. Producing the verdict.

10. user (2026-07-02T03:04:11.516Z)

Structured output provided successfully