

claude-windows / 6e9d0958-a2ec-4a8e-bb2d-6cce9833c527

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\6e9d0958-a2ec-4a8e-bb2d-6cce9833c527\subagents\agent-a430137c98f9dbf7c.jsonl`  
- SHA-256: `df2cbe578b5a7f6b54576a668a6184cc99ffcdd3224c0a1eb7d54f913208a58e`  
- Source modified: `2026-07-03T16:41:51+00:00`  
- Imported at: `2026-07-05T16:48:26+00:00`  
- project: `subagents`  
- session_id: `6e9d0958-a2ec-4a8e-bb2d-6cce9833c527`
```

Transcript

1. user (2026-07-03T16:40:26.331Z)

You are a code-review finder agent. Review target: PR #451 of the badminton-highlight-indexer repo, checked out at C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer on branch codex/alpha-readiness-plan (HEAD cd07339). Get the diff with: git diff origin/master..HEAD. Docs-only PR: new docs/ALPHA_LAUNCH_READINESS.md plus index additions.

ANGLE ? altitude. Check the change is implemented at the right depth given the repo's documented doc-lifecycle conventions. Read docs/DOC_STATUS.md (its lifecycle rules and §3a "tracked projects" conventions), docs/README.md structure, and any docs/archives/ conventions. Questions: (1) Does the repo's convention require anything the new tracked doc is missing (e.g., an ADR for a decision like D1-D5 "Decisions locked" ? check docs/archives/decisions/DECISIONS.md for how decisions are normally recorded; is locking new decisions D1-D5 inside a DRAFT tracker doc bypassing the ADR mechanism)? (2) Is a cross-repo tracker placed in one repo's docs/ the right altitude, given the repo set has sibling repos with their own NEXT_STEPS ? does an existing convention say where cross-repo trackers live? (3) The doc's §7 says it is "the source of truth" while NEXT_STEPS.md P0 item remains live ? does this create two competing sources of truth for the same P0 work, and does DOC_STATUS.md convention address that? Flag only concrete convention conflicts you can cite.

Return up to 6 candidate findings as a JSON array of {file, line, summary, failure_scenario}. Return [] if none. Your final message must be only the JSON array.

2. assistant (2026-07-03T16:40:28.457Z)

I'll analyze PR #451 for altitude issues related to the repo's documentation conventions. Let me start by examining the relevant files.

3. user (2026-07-03T16:40:30.006Z)

```
docs/ALPHA_LAUNCH_READINESS.md  
docs/DOC_STATUS.md  
docs/NEXT_STEPS.md  
docs/README.md
```

4. assistant (2026-07-03T16:40:31.754Z)

Now let me read the key documentation files to understand the repo's conventions:

5. user (2026-07-03T16:40:32.435Z)

```
1      # Documentation Status & Archival Register  
2  
3      > **Status:** ? **LIVING** ? the single board for doc **lifecycle + archival due-  
diligence**. **Last full review:** 2026-07-02.  
4      > **Update cadence:** whenever a doc changes lifecycle state (created / shipped a
```

milestone / delivered / archived), and re-review at each clean-slate.

5 > **Why this exists:** so we can (a) **review completeness** at a glance, (b) **update** it as work finishes to prove we're done, and (c) **never archive** a foundational doc without honestly updating it first (owner directive, 2026-06-21).

6

7 ---

8

9 ## 0. How to use this

10 - **Reviewing now?** Read §4 (the open findings/actions) ? that's what needs attention. §3 is the full register.

11 - **Marking something complete?** Update the doc's own **§7 progress tracker** (the source of truth for *its* completeness), then flip its row here, then ? if terminal ? run the §2 archival checklist.

12 - **This register does NOT duplicate completeness detail** ? it indexes lifecycle + archive-readiness and points to each tracked doc's §7.

13

14 ## 1. Lifecycle legend

15 `DRAFT ? IN PROGRESS ? DONE/DELIVERED ? ARCHIVED`. Plus: **LIVING** (permanent reference, never "done"), **PROPOSAL/DESIGN** (owner-gated, not started), **REPORT** (a terminal record of a run/launch), **CONVENTION** (a standing rule/guide).

16

17 ## 2. Archival due-diligence checklist ? *the repeatable process* (do NOT just `git mv`)*

18 Before a foundational/tracked doc moves to `docs/archives/`:

19 1. **VERIFY** claims against the code ? no stale assertions; tag load-bearing facts `[verified]`. (A doc that says "nothing built" when code shipped is *not* archive-ready ? fix it first.)

20 2. **Flip Status** ? DELIVERED / DONE / REJECTED + a completion note ? what shipped, where (PRs, paths, artifacts), and what was **deferred** / carried forward.

21 3. **Update ALL** inbound references ? `docs/README.md` index, `CLAUDE.md`, `docs/NEXT\_STEPS.md`, and any cross-doc links ? the new path + status.

22 4. **Keep the lineage** ? ADR forward/back-links + "supersedes / superseded-by" so the evolution stays traceable.

23 5. **git mv** the whole dir/file ? `docs/archives/past\_projects/` (terminal-state ONLY; live docs stay at `docs/` top level). Relative links inside become historical (note this in the doc's header).

24 6. **Record the move** ? flip the row here to §3e + a line in memory `current-state`.

25

26 > **Worked example** (the template): **DECOUPLED_COMPUTE** ? #270 ? verified M0?M2+M3.5 delivered on GCP, flipped ??? DELIVERED with a completion note (incl. carried-forward WASB-C3/DPA/md5), reindexed `docs/README` + `past\_projects/README`, kept the ADR-010?014 lineage, then `git mv` to `docs/archives/past\_projects/DECOUPLED\_COMPUTE/`.

27

28 ## 3. The register

29

30 ### 3a. Tracked projects ? lifecycle docs (archive when DONE; each doc's §7 is the completeness source-of-truth)

Doc	Status	Archive-ready?	Next action
`ALPHA_LAUNCH_READINESS.md`	DRAFT ? cross-repo alpha readiness tracker created 2026-07-03 after the golden corpus reached n=15; separates hosted product-alpha work from owned-model/commercial rally-detection gates	No	Execute §7 in order: corpus portability + n=15 baselines, hosted alpha serving, ablation/fusion reruns, then launch GO/NO-GO
`COMPUTE_DECOUPLED_SERVING/`	IN PROGRESS ? M0 approved + M1?M9 LARGELY SHIPPED (#279?#290; M5 #286, M6 #287, M8 #288, M9-auth #290 all merged, default-OFF); Gen-0 weights pushed (<code>weights/0.1.0-l4/</code>)	No	M11 flywheel + M12 gdrive-api (owner residual: counsel/prices/OAuth/real-run within the D17 cap)
`RALLY_QUALITY_RESEARCH.md`	Foundation, tracked & ACTIVE (§7)	No	Drives the R-series; complete when §7.1 single quality number is attributed
`R2_EVAL_METRIC_DESIGN.md`	IMPLEMENTED (augmenting)	No	(gates R4 + dual-eval; ablation pending) Promote-or-reject tolF1 as the gate (ablation) ? then archive with the R-cluster
`R1_MILESTONE_LAUNCH_REPORT.md`	LAUNCHED (#204) ? terminal record	Soon	(with the R-cluster) Keep with R-series for narrative; archive when the track wraps
`RALLY_DETECTION_GAP_CLOSURE.md`	DRAFT (owner-gated)	No	G1?G3 levers; part of

the R/quality cluster |

39 | `RALLY\_DETECTION\_FIXES\_PLAN.md` | **DONE** ? findings A/B/C resolved; P2b promoted
`inactivity\_temporal\_density` default-ON (#396) | With the R-cluster (deferred per its header;
gate = R2 toIF1 ablation + \$7.1 number) | Archive with the R-series when the track wraps |

40 | `RALLY\_DETECTION\_CODE\_REVIEW.md` | Terminal snapshot ? findings A (#394), B (#394 ?
promoted #396), C (`sdc#50`) all resolved | With the R-cluster | Archive with the R-series |

41 | `PER\_VIDEO\_WORKSPACE.md` | DRAFT, IN PROGRESS (P0 = #264) | No | ? supersede-reconcile
(finding #1) DONE; P3 de-hardcode base ? via #282; remaining P1?P6 (P1 v6 migration coordinates
with M8's v6 cost-ledger ? version bump authored once) |

42 | `MULTI\_SIGNAL\_FUSION\_PLAN.md` | P1 DONE; P2+ owner-gated | No | P3/P4 (learned fusion,
audio) ? needs GPU golden features |

43 | `EXPERIMENT\_HARNESS.md` | P1 IMPLEMENTED | No | A/B + promotion gate; supports the
quality track |

44 | `PROMINENT\_COURT\_DETECTION.md` | DRAFT (owner-gated) | No | C-series; multi-track
selection design |

45 | `EXECUTION\_POLICY.md` | P1 SHIPPED; P3b next | No | Async job policy; reused by Cloud
Run dispatch |

46 | `GOLDEN\_REGRESSION\_FIXTURES.md` | IN PROGRESS (owner-gated) | No | Golden eval
fixtures |

47 | **Data layer** ? `docs/data\_pipeline/` (`GOLDEN\_SET\_IMPLEMENTATION` .
`GOLDEN\_SET\_PHASE1\_VIDEOS` . `GOLDEN\_VIDEOS` . `VIDEO\_RECORDING\_GUIDELINES` .
`GOLDEN\_DATA\_SHARING`) | LIVING (data, **no-ML**) | n/a | Relocated 2026-06-21 to the isolated
data pillar; tracked via `data\_pipeline/README.md`. Roora *vendor* docs archived ?
`archives/roora-vendor/`. |

48 | `archives/past_projects/AUDIT_REPORT_khelsutra-
guru_2026-07-02-COMPLETED-2026-07-02.md` ** (umbrella) ** | ** ? COMPLETE / ARCHIVED (2026-07-02) **
? the five-repo code audit + folded-in *Remediation tracker* (P0?P28 all closed) + absorbed
prior audits. All non-owner-gated residuals landed; owner-gated/future items filed as issues
(#443/#444/#445, + #301/#327/#332/#384). | Archived | Owner confirmed archival 2026-07-02; moved
to `archives/past\_projects/` per \$2. Residuals now live as GitHub issues, not in this doc. |

49 | `DECISION\_SNAPSHOT\_REGRESSION.md` | **DRAFT ? decisions LOCKED** (D1?D11 locked
2026-06-25, #383). Zero code shipped for P1?P7. | No | Start P1+P2 (schema + stitch-plan
factor); de-risk the P10 main.py split |

50 | `DESKTOP\_APP\_PLAN.md` | DRAFT (owner-gated, nothing built) | No | P0 compute-
feasibility spike; relates to truly-local profile |

51 | `OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md` | PLAN/roadmap (owner-gated) | No (authoritative
roadmap) | Gen-0?Gen-2; the owned-model track |

52 | `PERSONALIZATION\_PLAN.md` | DESIGN ? owner-adopted; Phase-0 | No | Per-user models (DB
v5 store landed) |

53 | `UI\_ALPHA\_EXECUTION\_PLAN.md` | ACTIVE | No | khelsutra UI alpha track |

54 | `I18N\_PLAN.md` | Design (P0); **in implementation** (adapted) | No | Executed in
`khelsutra/docs/FRIENDLY\_AND\_MULTILINGUAL.md` (P1?site/, P3?web/); P2 backend stays here, owner-
gated |

55 | `POST\_166\_...RISK\_ASSESSMENT.md` | Plan for owner review | No | Milestones/risk doc |

56

57 ### 3b. Foundational / strategy references (LIVING ? archive only if **superseded**, not
when "done")

58 | Doc | Status | Note |

59 | ---|---|---|

60 | `PLATFORM\_ARCHITECTURE.md` | Strategy/research | The `ComputeBackend`/`StorageBackend`
registries ? realized by COMPUTE_DECOUPLED_SERVING |

61 | `COMMERCIALIZATION.md` | ? LIVING | Strategy; carries WASB-C3 |

62 | `VIDEO\_LOCALITY\_MODEL.md` | THINKING DOC | L0 fully-local = the truly-local profile |

63 | `HOSTING\_PLAN.md` | PROPOSAL | ** ? partly overtaken by ADR-013 + the shipped site ?
verify (finding #5) ** |

64 | `UI\_PROPOSAL.md` | PROPOSAL (partly shipped) | UX half **in implementation** ?
`khelsutra/docs/FRIENDLY\_AND\_MULTILINGUAL.md`; currency banner added (finding #5 ? for this doc)
|

65 | `STORAGE\_SHARING\_MODEL.md` | DRAFT (co-design pending) | ? |

66 | `ANALYZERS\_AND\_RECONCILERS.md` | DESIGN (owner-gated) | ? |

67 | `ROADMAP.md` / `CODE\_MAP.md` / `HOW\_RALLY\_DETECTION\_WORKS.md` /
`KHELSUTRA\_PRODUCT\_OVERVIEW.md` | LIVING references | Keep current; not archive targets |

68

69 ### 3c. Reports & terminal records (archive candidates once their track wraps)

70 | Doc | Status | Note |

71 |---|---|---|

72 | `REAL\_FOOTAGE\_VALIDATION\_REPORT.md` | validation complete | Motivates R2; archive with the R-cluster |

73 | `RALLY\_DETECTION\_QUALITY\_REPORT.md` | report | Quality snapshot; part of R-cluster |

74 | `QUALITY\_ITERATIONS.md` | LIVING empirical log | ****Keep**** ? the paper-backbone ablation log ([[paper-coauthor-aim]]) |

75 | `CODE\_AUDIT\_AND\_TEST\_HARDENING` · `HARDENING\_LOOP\_HANDOFF` · `DEFERRED\_HARDENING\_PLAN` | COMPLETED ? ****finish-archived 2026-07-02**** (top-level stubs removed; absorbed into the umbrella's ***Prior audits***) | ? Full copies under `docs/archives/\*-COMPLETED-2026-06-14.md`; inbound links repointed (finding #4 resolved) |

76

77 **### 3d. Conventions / guidelines / templates / runbooks (permanent ? do not archive)**

78 `PROJECT\_DOC\_TEMPLATE.md` · `NEXT\_STEPS.md` (? refreshed 2026-06-21) · `BACKLOG.md` · `EVALUATION.md` · `ABLATION\_LAYER.md` · `SETUP\_NEW\_MACHINE.md` · ****`CREATING\_REELS.md`**** (? end-user usage guide ? install ? process ? pick ? reel; LIGHT tier; added 2026-06-23) · ****`CONVERT\_GCS\_VIDEOS\_TO\_HIGHLIGHTS.md`**** (? operator runbook ? `gs://` videos ? reels on L4 Cloud Run; verified e2e 2026-06-23) · `TRACKNET\_WSL\_SETUP.md` · `README.md` · ***(2026-06-21 relocation: the data-layer guides moved ? `data\_pipeline/`; the field-ops + Roora docs ? `archives/{field-pmf,roora-vendor}/`)***

79

80 **### 3e. Already archived (`docs/archives/`)**

81 `past\_projects/DECOUPLED\_COMPUTE/` (? 2026-06-21, #270) · `past\_projects/low-regret-moves/` · ****`field-pmf/`**** (field-ops + customer-feedback, 2026-06-21) · ****`roora-vendor/`**** (annotation-vendor engagement, 2026-06-21) · `PER\_VIDEO\_OUTPUT\_LAYOUT-SUPERSEDED-2026-06-21.md` · `CODE\_AUDIT\_AND\_TEST\_HARDENING-COMPLETED-2026-06-14.md` · `DEFERRED\_HARDENING\_PLAN-COMPLETED-2026-06-14.md` · `HARDENING\_LOOP\_HANDOFF-COMPLETED-2026-06-14.md` · `COMMERCIAL\_READINESS\_REVIEW.md` · `MONETIZATION\_AUDIT.md` · `TEST\_RUN\_SUMMARY.md` · `research/` · `checkpoints/` · `quality-iterations/` ·

82

83 **### 3f. ADR log ? `docs/archives/decisions/DECISIONS.md`**

84 ****ADR-001 ? ADR-014**** (14, all ratified). The compute/serving lineage: ****ADR-005 ? 008 ? 009 ? 010 ? 011 ? 012 ? 013 ? 014****. ADR-014 was ****RATIFIED 2026-06-21**** (owner M0 sign-off for COMPUTE_DECOUPLED_SERVING).

85

86 **## 4. Open due-diligence findings / actions (2026-06-21 review)**

87

88 > ****? Doc-restructuring done (2026-06-21, owner-directed):**** field-ops/customer-feedback ? `archives/field-pmf/`, Roora annotation-vendor ? `archives/roora-vendor/`, and a no-ML ****data layer**** carved into `docs/data\_pipeline/`. Inbound + outbound links rewired; three dir READMEs added (the `data\_pipeline` one sets the "no segmentation/stitching" boundary).

89 1. ****? RESOLVED (2026-06-21) ? `PER\_VIDEO\_WORKSPACE.md` supersedes `PER\_VIDEO\_OUTPUT\_LAYOUT.md`.**** Folded forward the `video\_slug` naming option (workspace \$8.5) + the `output/GX010125\_run/` manual-prototype precedent; predecessor archived ? `docs/archives/PER\_VIDEO\_OUTPUT\_LAYOUT-SUPERSEDED-2026-06-21.md`; `CODE\_MAP.md` repointed.

90 2. ****? RESOLVED (2026-06-21) ? `NEXT\_STEPS.md` refreshed**** with a current 2026-06-21 state + a ****cross-track prioritized pending list****; the 2026-06-13 owned-model/quality body is kept (clearly marked) as that track's detail.

91 3. ****R-series / quality cluster is ACTIVE**** (`R1`, `R2`, `RALLY\_DETECTION\_GAP\_CLOSURE`, `RALLY\_QUALITY\_RESEARCH`, `QUALITY\_ITERATIONS`, `REAL\_FOOTAGE\_VALIDATION\_REPORT`, `RALLY\_DETECTION\_QUALITY\_REPORT`). ****Do NOT archive yet**** ? archive as a ***cluster*** once R2 is promoted-or-rejected as the gate + \$7.1's single number is attributed. (Owner already confirmed "leave as is", 2026-06-21.)

92 4. ****? RESOLVED (2026-07-02) ? finish-archived the hardening stubs.**** The three top-level pointer-stubs (`CODE\_AUDIT\_AND\_TEST\_HARDENING`, `HARDENING\_LOOP\_HANDOFF`, `DEFERRED\_HARDENING\_PLAN`) were removed; full archived copies remain under `docs/archives/`, inbound links (README, PLATFORM_ARCHITECTURE, DESKTOP_APP_PLAN) were repointed, and the sweep is summarized in the umbrella's ***Prior audits (absorbed)***.

93 5. ****`HOSTING\_PLAN.md` / `UI\_PROPOSAL.md` are PROPOSAL but partly overtaken**** by ADR-013 (DO dropped) + the shipped khelsutra site. ****? `UI\_PROPOSAL.md` banner added (2026-06-21)**** ? its UX half is now tracked in `khelsutra/docs/FRIENDLY\_AND\_MULTILINGUAL.md`; the `I18N\_PLAN.md` localization track is likewise in implementation there. ****`HOSTING\_PLAN.md` banner still pending**** (ADR-013 already noted in it).

94 6. ****? RESOLVED (2026-06-21) ? `COMPUTE\_DECOUPLED\_SERVING` M0 approved + M1?M4 SHIPPED****

(#279/#280/#282/#283); all §8 gating Qs locked (D7?D17); M5+ default-OFF code batch authorized (owner residual: counsel ToS/DPA, real GCP prices, OAuth app, real cloud run within the D17 \$100/?10k cap).

95

96 ## 5. Maintenance

97 Update a row when its doc changes lifecycle; run §2 before any archive; re-review §4 at each clean-slate. This register + the §2 checklist are the operational form of the `CLAUDE.md` archive policy and [[working-agreement-tracked-project-docs]].

98

6. user (2026-07-03T16:40:32.668Z)

1 - RALLY_QUALITY_RESEARCH.md ? **Foundation doc** for closing the local-vs-Gemini rally-detection gap: the measured gap + the code-grounded over-merge diagnosis ("shuttle moving ? rally"), an inventory of the eval/experiment harness + its gaps, the strengthened evaluation framework (segmental Edit/F1@k that make over-segmentation first-class, golden-corpus + active learning, Gemini-agreement as a *shadow* signal), an **adversarially-verified, cited external-research synthesis** (TAL/TAS boundary heads, racket-sports rally-state cues, small-data, eval rigor), a **prioritized cheap?durable roadmap** with falsifiable success criteria, and (§7) a **tracked project plan** ? work-item checklist, sequencing, the single quality **number** attributed to the effort, and a **definition of done** (project completes when all tiers land + the number is recorded). The plan that the QUALITY_ITERATIONS.md log executes against; ties together EXPERIMENT_HARNESS / MULTI_SIGNAL_FUSION_PLAN / ANALYZERS_AND_RECONCILERS / OWNED_MODEL_IMPLEMENTATION_PLAN.

2 - RALLY_DETECTION_GAP_CLOSURE.md ? **Tracked execution doc** (2026-06-18) for closing the remaining rally-detection accuracy gap **per video**, before **launch**, driven by the growing golden corpus. First ground-truth clip (GX010141) reframed it: boundaries are tight (R2/R3/R4 holds) ? the gap is **detection**: (G1) recall (local & Gemini both miss ~half on hard clips, and *different* rallies ? complementary), (G2) precision (local over-fires), (G3) **rally-END over-extension** on the quick-return-after-point (shuttle still moving ? needs a *player-state* signal). Levers (owner-gated): player-kinematics + shuttle-in-hand / sustained-invisibility end-rule with reverse-timeline backtrack, *rally_gate* Gate A wiring, complementary fusion. **Peer** of `RALLY\_QUALITY\_RESEARCH.md` (execution tracker, not research); §7 is the source of truth.

3 - PROMINENT_COURT_DETECTION.md ? **Design + experiment** (2026-06-19, default-OFF, owner-gated flip) for the **multicourt G2 over-fire**: identify the **prominent (foreground) court** (the one covering the majority of the frame) and keep only rallies whose shuttle lives inside it ? re-defining the shuttle signal `(x,y,visible)` ? `(x,y,visible-AND-inside-prominent-court)` **before** windowing, so a background/adjacent-court shuttle never forms a rally. 2D floor polygon (MVP) ? **pseudo-3D play-volume** (refinement, so high clears aren't clipped). Born from the GX010142 rally-#1 background-court FP; reuses *fusion_features.point_in_polygon* + *FusionConfig.court_polygon*. §8 progress tracker is the source of truth.

4 - RALLY_DETECTION_QUALITY_REPORT.md ? **Technical-paper checkpoint** (2026-06-17) of rally-detection quality: the honest measured state (over-segmentation milestone tolerance-F1 0.091?0.323 at n=11, p=0.002; merge-rate 0.694?0.150; first real-footage ground-truth at-par with START at Gemini parity), the task-faithful R2 metric, a detailed Gemini comparison, next steps + expected wins, areas yet to be explored, and a publishability / path-to-submission verdict. Renders to a polished PDF; the externally-shareable "toward a paper" summary of `RALLY\_QUALITY\_RESEARCH.md` + `QUALITY\_ITERATIONS.md`.

5 - EVALUATION.md ? How we measure and improve quality.

6 - EXPERIMENT_HARNESS.md ? Design (P1 SHIPPED) for A/B + shadow + SxS experimentation on rally-detection variants with **zero production regression** ? experiments are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift, rollback = flip a flag). Composes the existing *backend/eval/* machinery (*calibration*, *regression*, *metrics*); implemented in *backend/eval/experiment.py*.

7 - QUALITY_ITERATIONS.md ? **Living log** of rally-detection quality iterations (hypothesis ? what shipped ? measured/expected effect ? lesson). The reverse-chronological record; terminal-state snapshots get archived under *archives/quality-iterations/*.

8 - ANALYZERS_AND_RECONCILERS.md ? Design (owner-gated, NOT started): a **signal-fusion framework** for rally detection ? producers emit confidence-scored **signals** at three temporal **scopes** (frame detectors / window analyzers / session analyzers), a learned **scorer** weights them, and an auditable **report-first reconciliation**

pass lints the decision against the evidence. The spine: **no special-casing in the decision path** (signals?scorer, never `if/else`). Worked examples: a service-fault analyzer, a warm-up/practice span detector, and consuming a per-frame shuttle-state detector. A peer to `EXPERIMENT_HARNESS.md` and `MULTI_SIGNAL_FUSION_PLAN.md` (ADR-007 modular rally layer).

9 - PERSONALIZATION_PLAN.md ? Design (owner-adopted **hybrid**; Phase-0 landing) for the **per-user personalization** feature on a **generalization-first baseline**: a thin per-user **tuning** layer gated by the honest small-N statistics (a **min-N promotion floor** + a **structural-guard exemption**, never a service gate) + a **contribution flywheel** (free-tier videos pool into the owner-trained baseline; **"a tool that gets better the more you help it"**). Records the **locked owner decisions** (identity, storage, `min_n` semantics, Gate-A default, free-tier pooling). Built so far: per-user promotion gate (#141) + held-out-USER learning curve (#142).

10

11 - GOLDEN_REGRESSION_FIXTURES.md ? **Design** (owner-gated, NOT started): **video-free**, non-attributable regression fixtures so a clean clone runs the rally-seg suite with **no video/GPU/DB**. Two layers ? the post-detector freeze + **dual gates** (characterization + quality-floor) reusing the `backend/eval/` stack, and the **SEALED-FIXTURES** privacy filter (de-attribution-at-source, tiered public/key-gated, shuttle-only public tier). Carries a cited **IN/US/EU legal memo** (derived-data IP/privacy), a threat model, and a **\$7 progress tracker + definition-of-done**. Extends [GOLDEN_DATA_SHARING.md](data_pipeline/GOLDEN_DATA_SHARING.md).

12

13 **## ? Documentation status & governance**

14 - [AUDIT_REPORT_khelsutra-guru_2026-07-02 (? COMPLETE / ARCHIVED)](archives/past_projects/AUDIT_REPORT_khelsutra-guru_2026-07-02-COMPLETED-2026-07-02.md) ? the five-repo `khelsutra-guru` workspace code audit + its folded-in **Remediation tracker** (P0?P28 all closed). **Archived 2026-07-02** once every remediation wave landed and the last residuals were closed or filed as issues (#443 CORS, #444 corpus Phase-2, #445 GCS eval; + #301/#327/#332/#384). Absorbed the prior `CODE_CLEANUP_AUDIT_2026-06-24` + June-2026 hardening sweep (both archived).

15 - DOC_STATUS.md ? **living Documentation Status & Archival Register**: every foundational/tracked doc's lifecycle + archive-readiness, the **archival due-diligence checklist** (**"honestly update ? then archive"**), and the open doc-hygiene findings. Review here; update rows as docs complete.

16

17 **## Product / platform plans**

18 - ALPHA_LAUNCH_READINESS.md ? **Active tracked project** (`DRAFT`, created 2026-07-03): **cross-repo alpha readiness tracker** after the golden corpus reached 15 videos. Separates the hosted product-alpha lane (shareable upload?process?pick?reel) from the owned-model/commercial rally-detection gate (ship target + WASB-C3 still unresolved). \$7 is the source of truth for the concrete next PR sequence: corpus portability, n=15 baseline refresh, ablation/fusion reruns, hosted serving, alpha copy/launch review.

19 - [COMPUTE_DECOUPLED_SERVING/](COMPUTE_DECOUPLED_SERVING/README.md) ? **Active tracked project** (2026-06-21, `DRAFT` owner-gated): **the productionization successor** to DECOUPLED_COMPUTE (its deferred M3/M4). **One pure core** (DETECT?WINDOW?STITCH) that never branches on deployment profile, with a `ComputeBackend` (where it runs) + `StorageBackend` (what bytes) selected by **one startup config**. Ships **truly-local** (local/USB + local GPU, in-process stitch, zero cloud) + **cloud-serving** (upload<2GB / gdrive / bucket ? a **Cloud Run GPU job runs detect AND stitch**, metered into a per-job cost ledger) + **cpu-mock** (CI). Includes [TESTING_STRATEGY.md](COMPUTE_DECOUPLED_SERVING/TESTING_STRATEGY.md) (profile-PARITY proof for \$0), [architecture.svg](COMPUTE_DECOUPLED_SERVING/architecture.svg), a [runlogs/](COMPUTE_DECOUPLED_SERVING/runlogs/) posterity dir, and **ADR-014** ([lineage ADR-009?010?011?012?013?014](archives/decisions/DECISIONS.md)). \$7 M0?M10 tracker is the source of truth.

20 - [DECOUPLED_COMPUTE/](archives/past_projects/DECOUPLED_COMPUTE/README.md) ? **DELIVERED + ARCHIVED** (2026-06-21). **Lifted the rally pipeline** onto a **rented cloud-native GPU** doing bucket-driven `train` + `infer`: storage/compute seams (PRs #246?#259) + the **M3.5 native WASB runner**; **M0?M2 + M3.5 proven** on a GCP L4 (M2 ? `weights/0.1.0-l4/`). Delivered on **GCP** ? DigitalOcean dropped for compute ([ADR-013](archives/decisions/DECISIONS.md)). **M3/M4** (serving endpoint + per-video billing + scale-to-zero) deferred ? the "Cloud Run + GPU serving" subproject ([ADR-012](archives/decisions/DECISIONS.md)). Snapshot under `docs/archives/past_projects/DECOUPLED_COMPUTE/` (incl. `DEPLOY_REPORT_GCP_2026-06-20.md`).

21 - PACKAGING_PLAN.md ? **Active tracked project** (`IN PROGRESS`, 11 PRs merged 2026-06-21/22): **ship the engine** as a **downloadable batteries-included Python app** ? `pip install` ? `indexer --app` boots a local webserver and opens the **bundled KhelSutra**

SPA** single-origin, **torch-free** (LIGHT tier), cross-platform, no native component (decision D10/01). Built ON the decoupled-compute seam. **Supersedes the `DESKTOP\_APP\_PLAN.md` narrative.** \$7 tracker = source of truth (P0a?e ?, P1a/b ?, P2-launcher/P2a/P2d/P2e ?). **? Next:** P3 GPU validation + screencast ? P2b-install ? P2c wizard (cross-repo) ? P4 (ONNX/train/annotate = north star). Memory `[[packaging-app-plan]]` / `[[next-session-packaging-app]]`.

22 - DESKTOP_APP_PLAN.md ? **? Superseded narrative (by `PACKAGING\_PLAN.md`; D10 = self-hosted webserver, NOT a desktop app).** Original scoping doc (2026-06-18, `[design]`): package the local pipeline as a **cross-platform desktop app** for non-technical players ? plug in the camera's USB/SD card, click once, get back **only the highlight reels**, fully on-device (no upload, no original-video bloat). The centerpiece is the **de-WSL native-compute port** (run WASB natively per-OS: Linux/Windows CUDA · macOS MPS/CPU ? drop WSL2), and its **make-or-break unknown**: is CPU acceptable on a typical *no-GPU* enthusiast laptop (WASB is ~3.4× real-time on a laptop 2070S)? \$7 tracker = **P0 compute-feasibility spike ? P1 native detector ? P2 app shell ? P3 packaging/installers**. Realizes the **L0 fully-local** tier of VIDEO_LOCALITY_MODEL.md; is the **`LocalDesktop` ComputeBackend** of PLATFORM_ARCHITECTURE.md §3.2; the smaller/export-clean owned model (OWNED_MODEL_IMPLEMENTATION_PLAN.md) is one answer to the compute crux.

23
24 ## Practical / ongoing
25 - CREATING_REELS.md ? **? End-user guide:** install the LIGHT-tier app ? `indexer --app` ? process a video ? pick rallies ? download a watermarked reel. Task-first usage walkthrough (web UI + the API chain), with a \$6 pointer to the cloud compute option. LIGHT tier only (offline `motion` / optional Gemini; no on-device GPU or one-click cloud). Install mechanics ? root [README.md](../README.md); packaging roadmap ? PACKAGING_PLAN.md.

26 - CONVERT_GCS_VIDEOS_TO_HIGHLIGHTS.md ? **? Operator runbook:** `gs://` bucket videos ? highlight reels on the approved **L4 Cloud Run** (scale-to-zero, \$0 idle). Copy-pasteable: build image ? create the GPU Job ? cloud-serving control plane ? process(cloud)?query?compile?download ? prove the path (`verify\_compute\_path.py`) ? \$0 teardown. **Verified e2e 2026-06-23.** Raw deploy knobs ? ../deploy/cloudrun/README.md.

27 - SETUP_NEW_MACHINE.md ? **? The canonical "spin up a box ? ready for training & inference" runbook.** Copy-pasteable, human-followable, **verified end-to-end on DigitalOcean (2026-06-20)** ? but **compute is now GCP-only (ADR-013)**, see [deploy/gcp/README.md](../deploy/gcp/README.md): credentials ? provision ? bootstrap ? suite (922 passed) ? secret-free CPU inference (no GPU/keys) ? real Gemini inference ? Gen-0 training, with pointers to the GPU/native-detector + bucket paths. The `deploy/digitalocean/` scripts are provider-agnostic and reused on GCP.

28 - DATA_IN_GCS.md ? **? The corpus source-of-truth map: where every piece of data lives in `gs://kheIsutra-rally-corpus`** (canonical layout · current contents + drift · how `backend.worker` consumes it · the still-local-only **gap** · migration plan). Read this so a session pulls data from the bucket, not from someone's `C:\`/`F:\`/Drive ? the goal is to remove the local box as a blocker. Pairs with the two runbooks below.

29 - [deploy/gcp/nightly_regression/README.md](../deploy/gcp/nightly_regression/README.md) ? **? Nightly engine-regression OPERATIONS MANUAL (one-stop shop):** enable · run on demand · change any setting (email recipients/frequency/clips/config) · observe · pause/disable/remove. The full WASB-pipeline reel-count+cost+email nightly on an ephemeral GCP GPU VM (design/status: NIGHTLY_ENGINE_REGRESSION.md). The lighter CPU cached-trajectory tripwire is `scripts/nightly\_regression.py` (local Windows Scheduled Task, PR #202).

30 - TRACKNET_WSL_SETUP.md
31 - Evaluation harness and calibration drivers (in `backend/eval/`)

32
33 ## Data layer (collection · annotation · golden sets ? **no ML**)
34 - [data_pipeline/](data_pipeline/README.md) ? **the isolated data pillar:** how footage is recorded (`VIDEO\_RECORDING\_GUIDELINES`), golden sets are created + which videos (`GOLDEN\_SET\_IMPLEMENTATION` · `GOLDEN\_SET\_PHASE1\_VIDEOS` · `GOLDEN\_VIDEOS`), and data is shared/ingested (`GOLDEN\_DATA\_SHARING`). Deliberately **no segmentation/stitching/model** detail ? that's the ML pillar (top level + `COMPUTE\_DECOUPLED\_SERVING/`). The **annotation-vendor (Roora)** engagement is archived under `archives/roora-vendor/`.

35
36 ## Conventions & templates

```

37 - [PROJECT_DOC_TEMPLATE.md](PROJECT_DOC_TEMPLATE.md) ? Reusable skeleton for a **tracked
project doc**: status header · decisions-locked · progress tracker (??/?) · definition-of-done
· archived on completion. Standardizes the emergent pattern in `RALLY_QUALITY_RESEARCH.md` §7
and `GOLDEN_REGRESSION_FIXTURES.md`. Use for any substantial design/project work.
38
39 ## Historical documents
40 Completed plans, old research, ADRs, and checkpoints live in [archives/](archives/).
41
42 - **June-2026 test-hardening sweep** ? **COMPLETED** (2026-06-14; T1?T5 + full loop +
verification per #162) and **archived** ; the former top-level pointer-stubs were removed
2026-07-02 (absorbed into the umbrella's *Prior audits* (absorbed)). Full records:
[archives/CODE_AUDIT_AND_TEST_HARDENING-
COMPLETED-2026-06-14.md](archives/CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md),
[archives/DEFERRED_HARDENING_PLAN-COMPLETED-2026-06-14.md](archives/DEFERRED_HARDENING_PLAN-
COMPLETED-2026-06-14.md), [archives/HARDENING_LOOP_HANDOFF-
COMPLETED-2026-06-14.md](archives/HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md).
43 - **CODE_CLEANUP_AUDIT (2026-06-24)** ? the prior deep engine audit (35 findings; Phase
1 complete, Phase 2 largely shipped, remaining = the `main.py` god-module split). **Absorbed
into the umbrella** (see its *Prior audits (absorbed)*) and archived ? [archives/past_projects/
CODE_CLEANUP_AUDIT_2026-06-24-SUPERSEDED-2026-07-
02.md](archives/past_projects/CODE_CLEANUP_AUDIT_2026-06-24-SUPERSEDED-2026-07-02.md).
44 - **KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN** ? the five-repo audit execution board,
**folded into the umbrella** as its *Remediation tracker (live)* and archived ?
[archives/past_projects/KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN-
FOLDED-2026-07-02.md](archives/past_projects/KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN-
FOLDED-2026-07-02.md).
45 - [DECISION_SNAPSHOT_REGRESSION.md](DECISION_SNAPSHOT_REGRESSION.md) ? **? Tracked
project (decisions D1?D11 LOCKED 2026-06-25, #383; implementation P1?P7 pending):** a
deterministic **quality-neutral proof** ? freeze + compare the pipeline's **decision output**
(rally segments + stitch/concat plan) off a cached trajectory, GPU-free, in CI. The behavioral
proof that lets a *genuine* (non-AST-identical) refactor be confidently called quality-neutral.
**Never hashes the encoded reel** (non-deterministic). Extends the existing characterization
snapshot (`test_golden_fixtures.py`) from segmentation to the stitch plan; peer of the M4
quality-floor + GPU `nightly_reelcount.py`. §7 tracker is the source of truth.
46
47 - **Field-ops / customer-feedback** (recruiting flyer · field-visit playbook · PMF) ?
[archives/field-pmf/](archives/field-pmf/) (archived 2026-06-21; PMF validation still an
active *business* thread ? see `NEXT_STEPS.md`).
48 - **Roora annotation-vendor** engagement (labeling brief · vendor bake-off) ?
[archives/roora-vendor/](archives/roora-vendor/) (archived 2026-06-21).
49 - Superseded: per-video output layout ? [archives/PER_VIDEO_OUTPUT_LAYOUT-
SUPERSEDED-2026-06-21.md](archives/PER_VIDEO_OUTPUT_LAYOUT-SUPERSEDED-2026-06-21.md) ( ?
`PER_VIDEO_WORKSPACE.md`).
50
51 See `archives/README.md` for the umbrella policy and `archives/past_projects/README.md`
for classification of small tech debt / starter projects.
52

```

7. user (2026-07-03T16:40:32.871Z)

```

1 # Next steps (live) ? refreshed 2026-06-21
2
3 > **How to read:** this top section is the **current cross-track state + prioritized
pending list** (refreshed 2026-06-21, replacing the stale 2026-06-13 header). The detailed
**owned-model & rally-quality track** plan (2026-06-13) is **preserved below, still valid** ? it
is now *one of several* active tracks. Live blow-by-blow: memory `current-state.md`. Doc
lifecycle/hygiene: `docs/DOC_STATUS.md`.
4
5 ## ? Where we are (2026-06-21)
6 **DECOUPLED_COMPUTE delivered on GCP + archived** (#270); **compute = GCP only**
(ADR-013, DO dropped). **`COMPUTE_DECOUPLED_SERVING`** subproject spun up (#272, **ADR-014**) ?
one pure core + local/cloud profiles; **M0?M9 LARGELY SHIPPED** (ADR-014 ratified 2026-06-21;
M1?M4 #279?#283, M5 #286, M6 #287, M8 #288, M9-auth #290 all merged, default-OFF). **Per-video
workspace P0** shipped (#264); **doc-status register + archival due-diligence** in place (#273).
The **khelsutra site + per-rally UI** ship in parallel (sibling track). The **rally-quality /

```


owned-model** track (below) is the core ML work, gated on **growing the golden corpus**.

7

8 > **? LAUNCH DECISION ? NO-GO (owner, 2026-06-30).** The rally-detection product is **not** cleared for an external/commercial launch; we stay in **development**. This is the expected default for the dev phase. Two unblockers gate a future GO: **(1)** set + meet the **ship-gate target (#172)** ? currently unset, needs corpus growth (and the R2 toLF1 ablation, item 8); **(2)** resolve the **WASB shuttle-detector weight provenance** (commercial-clean licensing). P2b (#396) was a do-no-harm quality increment, **not** a launch gate ? merging dev PRs ? launching. Re-decide only when both unblockers clear.

9

10 ## ? Prioritized pending ? all tracks

11 **P0 ? highest leverage / unblocks the most**

12 0. **? [owner-priority, 2026-06-21] MAKE IT UI-DRIVEN + ALPHA-SHAREABLE ? "a few recordings lying around ? highlights, from a page, no SSH".** Today the cloud flow is **SSH/CLI-only** (`docs/CLOUD\_GPU\_DRYRUN\_GUIDE.md`); **M12 does NOT fix it.** Needs (a separate khelsutra `web/` SPA + deploy track, NOT an engine seam): **(i)** deploy the engine as a reachable **service** (M7 ? Cloud Run service, ? owner GCP spend for a *persistent* endpoint), **(ii)** the **SPA ingest/progress screen** (khelsutra `web/` B-P2: enter/upload video ? `/api/process` ? poll `/api/jobs` ? reel; ZERO engine code), **(iii)** flip **M9 edge-auth ON** (merged #290, default-OFF) to gate alpha tenants, **(iv)** a **compute PICKER in the UI** ("your machine vs our cloud", D13/D15). The engine seams make it *possible*; the SPA makes it *usable*. Full record: [[ui-driven-cloud-serving-gap]]. **Concrete execution tracker after the n=15 corpus review:** `docs/ALPHA\_LAUNCH\_READINESS.md` §7.

13 0b. **? [packaging track, IN PROGRESS ? 11 PRs merged 2026-06-21/22] Downloadable batteries-included app ? `docs/PACKAGING\_PLAN.md`. The LIGHT-tier engine is bundle-ready: `pip install` ? `indexer --app` serves the **bundled SPA** single-origin, **torch-free**, with app-home state / ffmpeg resolver / DNS-rebinding guard / single-instance lock / safe DB upgrades (P0a?e, P1a/b, P2-launcher/P2a/P2d/P2e). **? Next (fresh focused pushes):** **(1)** P3 **GPU validation + screencast** (L4, ~\$0.50, DELETE-after, \$100 cap ? `[[gcp-vm-always-delete]]` ? native-WASB 0-rally tuned-config blocker), **(2)** P2b-install (uv/script + truly-local default config), **(3)** P2c first-run wizard (cross-repo khelsutra SPA from `/api/capabilities`), then P3-cloud (Cloud Run) + P4 (ONNX/train/annotate = north star). Memory `[[packaging-app-plan]]` / `[[next-session-packaging-app]]`.

14 1. ? **`COMPUTE\_DECOUPLED\_SERVING` M0?M9 LARGELY SHIPPED** (2026-06-21, ADR-014 ratified). M1?M4 (#279/#280/#282/#283), **M5 #286, M6 #287, M8 #288, M9-auth #290** all merged (default-OFF); the **real M7 L4 run** validated the worker (22 rallies, VM deleted, \$0). All \$8 gating Qs locked (D7?D17). **Owner residual** (not blocking Claude's default-OFF code): counsel ToS/DPA (M9-consent live + M11 live), real GCP prices+FX (M8 calibrate), Google OAuth app (M12 real), CF team-domain/AUD + ingress-lock; real cloud verification authorized within a **\$100/?10k cap** (D17).

15 2. **[owner, ongoing] Grow the golden corpus** ? multi-court badminton + ?3 matches/sport (TT first). The single highest-leverage move; the quality track *and* full-corpus training feed on it.

16 3. **[owner + Claude . GPU] Full-corpus training (n?5)** ? a real `weights/<ver>/` (the `0.1.0-l4` bundle is n=2 plumbing-only).

17

18 **P1 ? near-term, mostly Claude-doable (default-OFF, green-gated)**

19 3b. **[Claude] Five-repo audit health remediation ? ? COMPLETE (2026-07-02).** All waves (P0?P28) landed and the last residuals were closed or filed as issues; the umbrella tracker was archived to `docs/archives/past\_projects/AUDIT\_REPORT\_khelsutra-guru\_2026-07-02-COMPLETED-2026-07-02.md`. Remaining owner-gated/future work lives as GitHub issues: #443 (CORS default), #444 (corpus Phase-2 upload), #445 (GCS eval plumbing), plus #301/#327/#332/#384.

20 4. **[Claude] `COMPUTE\_DECOUPLED\_SERVING` batch ? REMAINING: M11 + M12** (M5/M6/M8/M9-auth ? merged this session). **M11** = `backend/flywheel.py` capture?store?shadow-eval?goldenize plumbing reusing `workspace.py::for\_video` + `backend/eval/experiment.compare` + `backend/eval/promotion.py::promote\_if\_served\_safe`, behind `flywheel.\*`=OFF + a faked `consent\_attested` (? privacy-sensitive ? give it fresh focus). **M12** = `backend/storage/gdrive\_api.py` `GDriveApiStorage(StorageBackend)` + a `gdrive-api://` scheme (regex already allows the hyphen; add to `\_KNOWN` + `\_backend\_class` + the StorageConfig Literal) against an INJECTED fake Drive (mirror `gcs.py` DI); keep the mount-only `gdrive://` untouched. ? **M9-consent cols = a new v7 migration** (M8 took v6). M9-consent ToS/DPA text = owner/counsel.

21 5. **[Claude] Per-video workspace P1?P6** (P1 v6 DB migration ? P2 worker ? P3 per-video routing of `detector/+`tmp/` \*[[the `output/chunks_` de-hardcode base = ? done via #282]]* ? P4

? P5 Launch-Report flip ? P6 GPU parity). ? P1's v6 migration coordinates with the M8 v6 cost-ledger migration ? author the version bump once.

22 6. ? ****Ops-Agent codified (#291)**** ? `deploy/gcp/create\_gpu\_vm.sh` (the one-command VM helper) always bakes the Ops-Agent startup-script; provision.sh grants the APIs+SA roles. PROVEN active on a fresh VM 2026-06-21. ([[gcp-ops-agent-enable]])

23 7. ****[owner]** Set the ship-gate target (#172)** ? needs corpus growth.

24 8. ****[owner-side GPU]** R-series: promote-or-reject `R2` tolF1 as the gate** (the ablation) ? needs GPU-extracted golden features.

25

26 ****P2 ? later / gated / deferred****

27 9. Rally-quality ****P3/P4**** (person-fusion LOVO, audio cue) ? needs GPU golden features.

28 9b. ? ****DONE (#396):**** the R4 density fix (finding B / P2b in `RALLY\_DETECTION\_FIXES\_PLAN.md`) is ****promoted to default-ON****. Golden-set A/B at the ****SERVED**** windowing showed tolF1 0.353?0.364 (?tolF1 +0.010, no regression), so `inactivity\_temporal\_density` default was flipped to True. (Landed flag-gated default-OFF in #394.) Fixed the sparse-tracking quirk where one fast sample after a detection gap held the rally tail open.

29 10. ****M0.4 scorer swap**** ? TemporalMaxer (per the 2026-06-17 external review).

30 11. ****Multisport**** (table-tennis first, then pickleball).

31 12. Serving ****M8/M9**** ? per-video cost ledger, edge auth, DPA/consent (within `COMPUTE\_DECOUPLED\_SERVING`; owner/legal-gated).

32 13. ****WASB-C3**** weight provenance (gates *sharing* a trained model) ? legal.

33 14. ****Rename repo off "badminton"***** (owner, BACKLOG).

34 15. Owner housekeeping: retire the shared ****DO droplet**** + ****DO token****; collector ****md5**** export (collector repo).

35 16. ****PMF field validation**** (owner-side, business track).

36

37 ****Doc-hygiene** (tracked in `docs/DOC\_STATUS.md` §4):** ? #1 per-video reconcile + ? #2 this refresh; ****pending:**** field/PMF/Roora ? archives + a data-layer `docs/` dir (owner boundary confirm); confirm the COMPLETED top-level stubs; verify `HOSTING\_PLAN`/`UI\_PROPOSAL` currency.

38

39 ---

40

41 **## Owned-model & rally-quality track ? detail (2026-06-13 plan; still valid)**

42

43 ****Updated:**** 2026-06-13 (****rally-detection quality track****: fusion + experiment-harness design docs merged,

44 owner-gated ? see the dedicated section below, the desktop-agent pick-up). Prior: 2026-06-10 (audit?M0?Gemini-battery session; later same day the ****UI-alpha track opened****:

45 ADR-009 ratified ? KhelSutra Guru / khelsutra.guru, private `khelsutra` repo, local-first delivery ? see

46 `docs/UI\_ALPHA\_EXECUTION\_PLAN.md`; ****PR-1** async jobs (#16) MERGED 2026-06-11, PR #87**;

47 ****PR-2**

47 upload/download MERGED 2026-06-11, PR #99** ? live-E2E-verified (3-part chunked upload MD5 byte-identity +

48 compile?browser-download); next: PR-3 identity. Same day: the ****8?9 quality sweep #90?#98 landed**** ? ruff/mypy/

49 coverage-floor CI lanes, SQLite conn-close fix, config unknown-key warning, hybrid-twin collapse into

50 `trajectory\_hybrid.py`, WSL tilde+abspath live-path fixes; suite 388 green; LOVO 0.626 re-verified exact).

51 ****Authoritative roadmap:****

52 `docs/OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md` (+ ADR-008 for topology). ****Latest blow-by-blow:**** memory

53 `current-state.md` / `session-handoff.md`. This file = the prioritized "what to do next," refreshed each session.

54

55 **## ? Next-session / GPU-box pickup (2026-06-15)**

56 ****M0** in-container quick wins are COMPLETE:** ship-gate provisional blocker (****#170****), train/eval split guard (****#171****), the ****doc-consistency sweep**** (****#174**** ? corrected 16 drifted docs vs the code), and the ****P3 person-feature fusion litmus machinery** is already present + suite-green** (`backend/eval/fusion\_compare.py` + `ablation.py`: person-driven ?F1 / bootstrap CI / paired sign test; the real-golden `skipif` litmus waits only on GPU-extracted data). ? CI is red repo-wide due to a ****GitHub Actions minutes/billing outage**** (account-level, not code) ? work is local-gated + `--admin`-merged until Actions is restored (Settings ? Billing

? Actions minutes / spending limit).

57

58 **Immediate next step ? the real P3 LOVO measurement** (this box has a GPU +
`torch+cu124` + `cv2`; videos and trajectory CSVs are present, but the GPU-extracted golden
features + `.rallies.csv` labels are not yet here):

59 1. **[owner / GPU box]** `python -m backend.eval.fusion_golden --manifest
output/human_lovo_manifest.json --out-dir "%USERPROFILE%\OneDrive\rally-annotator\golden-
data\features\2026-06-15-n6" ? produces `\*\_golden\_features.json` for the 6 golden videos. (GPU
+ WSL + videos + trajectory CSVs + `.rallies.csv` labels are owner-side only.)

60 2. Artifacts land in the **shared cloud folder** `C:\Users\avidu\OneDrive\rally-
annotator\golden-data\` (OneDrive auto-syncs both ways). Convention + workflow:
[`GOLDEN\_DATA\_SHARING.md`](data_pipeline/GOLDEN_DATA_SHARING.md); tracked in **#175**.

61 3. **[CPU dev/agent session]** `python -m backend.eval.fusion_compare --features-dir
<shared>\features\2026-06-15-n6 --record` + `python -m backend.eval.ablation --features-dir
<shared>\... --granularity group` ? real person-feature ? F1; copy the JSONs into `output/` to
also pass the `skipif` litmus tests (`tests/test\_ablation.py::test\_litmus\_reproduces\_gate\_a`).

62

63 **Also queued (owner-gated / pending input):**

64 - Ship-gate target calibration ? **#172** (owner decision; needs corpus growth).

65 - Tests reorganization ? blocked on the owner's grouping metric (provide it ? an agent
will do it).

66 - Restore CI ? the Actions outage is account-level (billing/minutes), not a repo/config
bug.

67

68 **Leaving (bulky / M2 ? no start without owner go):** ActionFormer (M0.4), event-index
DB migration (M0.1), Gemini-silver purge (M0.3), cost-to-Gen-2 benchmark + ByteTrack
deprecation, multisport, repo rename, PMF execution.

69

70 ## Where we are (one paragraph)

71 ADR-008 ratified ("repo split driven by productionization, not isolation"): training
lives in-repo (`training/`)

72 behind a CI-enforced import wall; the featurizer is single-sourced
(`backend/features/rally\_seq.py`, train==serve);

73 the **Gen-0 harness ships** (`python -m training.gen0.harness --manifest
output/human_lovo_manifest.json --floor
eval_baselines/heuristic_lovo_n6.json --version <semver>`) and produced artifact v0.1.0
(LOVO **0.626**, floor passed,

75 sign test 5/6 wins $p=0.109$? honestly not significant, ship=False). The **n=6 owned
corpus is in the collector**

76 (262 rallies, Proprietary, commercial export = owned data only after the ShuttleSet
reclassification). The **Gemini

77 battery is done** (see table) ? the moat is quantified.

78

79 ## The quality table that now drives strategy (IoU 0.5 vs human golden labels)

80 | Path | Mahadevpura (clean-ish) | Kushagra (multi-court) |

81 |---|---|---|

82 | Heuristic (velocity windowing) | 0.657 | 0.157 |

83 | **Gen-0 owned head** (LOVO, n=6) | 0.744 | 0.561 |

84 | Two-stage WASB+Gemini refine | 0.83 | ? |

85 | **Gemini wrapper** (`gemini-flash-latest`) | **0.93** | **0.66** |

86

87 **Reading:** clean footage is commoditized (serve it with Gemini for cents). **Multi-
court drops the commodity

88 wrapper to 0.66 ? that is the owned model's ship target** (its FPs are background-game
pickups + dead-time merges,

89 the exact contrast-metric confound). Gen-0 is 0.10 behind with only n=6 ? a corpus-
growth-sized gap, not an

90 architecture-sized one. Two-stage refine = the cost-efficient serving path for LONG
tapes (uploads candidate clips

91 only); wrapper 503-flakiness (observed all session) makes the provider-fallback registry
load-bearing.

92 Baselines tracked in `eval\_baselines/` (heuristic_lovo_n6.json,
gemini_wrapper_2026-06-10.json).

93

94 ## Rally-detection quality track (design merged 2026-06-13, owner-gated; implementation

in progress)

95 **This is the track a desktop/local agent is slated to pick up.** Design docs landed (PR #112 + session). All implementation is **owner-gated** (per explicit in the plan); each step = ONE PR off `master`.

96

97 **Hardening loop (T1-T5 positive tests + audit + finale) COMPLETE (2026-06-14, PRs #156/#157/#160/#161/#163 + #165 + final records).** See archived `docs/archives/CODE\_AUDIT\_AND\_TEST\_HARDENING-COMPLETED-2026-06-14.md` (Living Log with every pre/post/review) and `docs/archives/HARDENING\_LOOP\_HANDOFF-COMPLETED-2026-06-14.md` (full rules + STATE). Top-level now pointer stubs per #165 review. All 5 items + verification + archive done; clean slate.

98

99 - **Docs:** `docs/MULTI\_SIGNAL\_FUSION\_PLAN.md` (shuttle + person + audio, one fusion layer) .

100 `docs/EXPERIMENT\_HARNESS.md` (A/B + shadow + promotion gate, zero-regression) .

101 `docs/ROORA\_LABELING\_GUIDELINES.md` (v8) . `docs/HOW\_RALLY\_DETECTION\_WORKS.md` (2-layer mental model).

102 - **Key finding (no re-investigation needed):** the owner's "held-shuttle counted as a rally" bug is **NOT a missing capability** ? `backend/pipeline/detectors/rally\_gate.py` already implements the net-crossing **"Gate A"** (`apply\_gate(..., min\_crossings=2)`) that kills service-feed/held-shuttle windows, but it is **only called from eval drivers** (and there was no `min\_crossings` knob in `IndexingConfig`). **Wiring Gate A into the live segmenter (now via fusion_hybrid) was the single cheapest, highest-confidence quality win.** (P2 FusionSegmenter specifics: registry `fusion\_hybrid`, **default-OFF**, seeds from substrate, persists `net\_crossings` + optional person features, optional served Gate A/B via config knobs; adversarially reviewed; suite green. See `tests/test\_served\_gate\_a\_regression.py` for the honest finding that on production windowing Gate A is neutral/negative in some cases, hence kept opt-in with knob for safety, and the leaderboard revert note.)

103 - **Current status (2026-06-14):** P1 (person cue extraction to reusable `PersonDetector`) + P2 (FusionSegmenter + Gate A signal + served handoff gating) + experiment harness P1 DONE (adversarial review + NO-REGRESSION; suite green). **P3 next (this PR + follow-ups):** learned fusion (person features ? harness `compare` LOVO lift on the 6 golden videos; eager-person-compute optimisation); then P4 audio via hit_detector.

104 - **Pick-up order (each = ONE PR off `master`, owner approval first, flag-gated default-OFF, promote only on measured LOVO):**

105 1. (P2 first-step, largely done via fusion_hybrid) Wire Gate A into production served path + config knob (min_crossings); apply in hybrids at runtime. (Now available opt-in; served neutral on production windowing per litmus ? kept opt-in with knob for safety.)

106 2. P3 first step (this work item): harness `compare` LOVO litmus for person-on variant (players_in_court + two_sided + colocation etc. from `fusion\_features`) on the 6 golden; add the eager compute optimisation note/landing. Pure + harness (testable here).

107 3. P3/P4 continuation: audio cue integration + full learned promotion if lift proven.

108 - **Discipline (from EXPERIMENT_HARNESS.md):** every new cue **flag-gated, default OFF**; promote to default **only on measured LOVO lift** through `run\_regression.py` (exit 0/1/3); pinned `CONTROL` ? rollback = config flip. Harness re-uses ~80% existing code.

109 - ? Gate A, person cue logic, harness, and pure fusion paths are testable here; the heavier real-video / golden LOVO / ablation confirmation waits on the GPU-extracted golden features + `.rallies.csv` labels, which are not yet present in this checkout.

110 - **External due-diligence review inputs (2026-06-17) ? PICK UP AFTER current ideas are exhausted; full details + citations in [`RALLY\_QUALITY\_RESEARCH.md`](RALLY_QUALITY_RESEARCH.md) §5.6.** An owner-commissioned 31-reference review of the quality report *validated* the methodology and surfaced these **future** levers (research inputs, not active work, not new claims; verify citations before any submission):

111 1. **"First-mile" paper framing (highest-value):** stroke-classifiers (BST/DSTA) + the big datasets (ShuttleSet ~36k strokes, Fine-Badminton ~27k actions) all **assume a pre-trimmed rally clip already exists** ? we solve the ignored prerequisite. Pitch the paper as **"the missing untrimmed-video preprocessing pipeline that lets BST/MonoTrack run on raw footage."*

112 2. **Ground R2 in Action-Spotting / Precise-Event-Spotting** (SoccerNet/TenniSet/T-DEED) ? turns R2 into a defensible publishable contribution.

113 3. **Add domain baselines for any submission:** MonoTrack + BST on ShuttleSet / Fine-Badminton (we have only heuristic + Gemini today).

114 4. **Prefer TemporalMaxer over ActionFormer** for the M0.4 scorer swap (parameter-free max-pool, ~2.8x fewer GMACs, small-N robust) ? see "M0.4 next increment" below.

115 5. **Court polygons ? pose ? the rally-END lever:** masking on-court players can resolve *occluded end-of-rally states* (the #1 remaining gap vs Gemini), not just spectator

```

noise.
116      6. **Name the eval?serve lesson "pipeline distribution skew" ? a citable paper
"lesson."
117
118      **Discipline reminder:** owner approval before starting any owner-gated item; one PR per
step; babysit the PR (poll/fix/reply) until merge observed, *then* advance.
119
120      ## Prioritized next steps
121      1. **[owner, in progress] GROW THE GOLDEN CORPUS** ? the single highest-leverage move;
every lever below feeds on it.
122      Priority order for new videos: **(a) multi-court badminton** (the target distribution
AND where the ship target
123      lives), **(b) ?3 matches per new sport ? table tennis first** (WASB transfers at F1
0.909; pickleball labels are
124      corpus-gold but not trainable until a clean MIT/Apache ball detector exists), (c)
capture variety per
125      `VIDEO_RECORDING_GUIDELINES.md`; **adult sessions only for the training pool**
(consent guardrail). Per video:
126      label ? `curate import-local --normalize --license Proprietary --fps <true fps>` ?
re-run the Gen-0 harness.
127      The paired sign test needs n: at n=10 with 9 wins, p?0.011 ? significance is
reachable this month.
128      2. **[owner decision] Set the ship-gate target** ? now informed: floor = heuristic 0.480
(necessary), **multi-court
129      reference = wrapper 0.66** (the number that makes the owned model worth serving),
clean-footage ceiling = 0.93
130      (not the target ? Gemini serves that tier). Suggested shape: "LOVO F1@IoU0.5 ? 0.70
on multi-court folds with
131      paired-sign p<0.05 vs heuristic" ? owner to ratify/adjust.
132      3. **M0.4 next increment ($0/local):** swap the Gen-0 LogReg for a lightweight temporal
head (MIT, minutes on the
133      2070S ? NOT the heavyweight Gen-2 stack) inside the existing harness; per-video
threshold calibration (the
134      LOVO-vs-oracle gap is visible per fold in the harness output). **Prefer
TemporalMaxer** (parameter-free
135      max-pool, ~2.8x fewer GMACs, small-N robust) over ActionFormer/ASFormer per the
2026-06-17 external review
136      (`RALLY_QUALITY_RESEARCH.md` $5.6) ? better fit for the cheap/local/fast mandate.
137      4. **Multisport thread:** ingest Roora pickleball/TT CSVs (`import-local --normalize`);
merged-prototype LOVO across
138      badminton+TT once TT has ?3 matches; identify a clean pickleball ball detector. ? TT-
on-WASB serving stays gated
139      by C3 (provenance multiplies per sport).
140      5. **PMF validation (cheaper than any engineering month):** C13 academy interviews (+
the two added questions:
141      "multi-court footage nobody watches?" / "pay per-match to index it?"); camera pilot
at the warmest 2-3 academies
142      (consent at capture; demo served via the Gemini path or human-polish ? the reel
pattern exists:
143      `output/MahadevpuraSingles_highlights.mp4`). **Field kit COMPLETE (2026-06-12, owner
ratified the
144      physical-visit approach; reframed product-feedback-first per owner spec):**
145      `docs/PRODUCT_FIELD_ASSOCIATE_FLYER.md` (the recruiting flyer ? owner fills ?/contact
placeholders and recruits) +
146      `docs/FIELD_VISIT_PLAYBOOK.md` (the hired associate's visit playbook: coach + student
script, demo-after-Q6
147      rule, ground rules, Phase-2 record-and-show-back stretch goal) +
`docs/FIELD_VISIT_ANSWER_SHEET.md` (per-visit
148      capture form incl. demo-reaction + student blocks) + `docs/PMF_FIELD_SIGNALS.md`
(G/A/R anchors + **pre-registered
149      decision rules** ? PROCEED / PIVOT A-C / FALSIFIED thresholds fixed before the first
visit).
150      6. **Housekeeping:** ~~rotate the Gemini API key~~ **DONE 2026-06-10** (owner-verified
working; confirm old-key
151      deletion in AI Studio); **rename repo off "badminton" (owner action,

```

```

152         in BACKLOG); collector=SoR confirmation; off-peak niceties: none pending ? the
battery completed.
153
154     ## Pointers / artifacts
155     - Roadmap: `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` · Topology: ADR-008
(`docs/archives/decisions/DECISIONS.md`) ·
156     Study: `docs/archives/research/OWNED_MODEL_TRAINING_STUDY.md` · Quality history:
157     `docs/archives/quality-iterations/2026-06-multivideo-lovo/`.
158     - Harness: `training/gen0/harness.py` (one command ? train ? LOVO ? gate ? artifact
bundle; weights gitignored,
159     manifest reviewable). Featurizer: `backend/features/rally_seq.py` (FEATURE_SCHEMA v1).
160     - Corpus: collector DB (system-of-record, 6 matches/262 rallies) · eval manifest
`output/human_lovo_manifest.json`
161     (gitignored) · durable trajectories + Roora files: `C:\Users\avidu\Projects\Annotation
Setup\Collect\Trajectories\` ·
162     labels: `...\Collect\Golden Labelled\`.
163     - Gemini ops learnings: serial uploads (`indexing.ai_max_workers=1`) + chunk re-encode
164     (`indexing.ai_chunk_scale_width=1280`) + 6-attempt exponential generate retry ? a
cold/prepaid project burst-throttles
165     with a MISLEADING "credits depleted" 429; ramp gradually.
166     - ? No owned-model implementation beyond greenlit M0.x until owner greenlights the next
milestone. Committed next
167     PLATFORM slice = async job model (single-tenant).
168

```

8. assistant (2026-07-03T16:40:35.375Z)

Now let me read the new ALPHA_LAUNCH_READINESS.md file:

9. user (2026-07-03T16:40:35.821Z)

```

1     # Alpha Launch Readiness Plan
2
3     > **Status:** `DRAFT` - cross-repo execution tracker; **Owner:** @avidullu; **Created:**
2026-07-03; **Last updated:** 2026-07-03
4     > **Lifecycle:** `DRAFT` -> IN PROGRESS -> DONE -> archived`
5     > **Tracking anchors:** Section 7 is the progress tracker and source of truth; indexed
in `docs/README.md` and `docs/DOC_STATUS.md`.
6     > **Relation to existing docs:** consolidates the launch-relevant pending items from
`NEXT_STEPS.md`, `ROADMAP.md`, `data_pipeline/GOLDEN_VIDEOS.md`, `R2_EVAL_METRIC_DESIGN.md`,
`COMPUTE_DECOUPLED_SERVING/`, and the sibling `khelsutra` Studio docs.
7     > **Honesty note:** `[verified]` means checked against the local synced repos on
2026-07-03. `[analysis]` means derived from a local command run on that snapshot. This is not
legal, financial, or launch approval.
8
9     ---
10
11     ## 0. TL;DR
12
13     The corpus is now large enough to move from "grow labels first" into alpha-readiness
work, but the 15-video headline is not itself a launch gate pass. `[verified]` The golden
manifest has 15 badminton videos; `[analysis]` the local manifest needed path rebasing before
manifest-driven commands ran; `[analysis]` only 6 of 15 feature JSONs had the full fusion
schema; `[analysis]` the n=15 quality baseline is stale and must be intentionally refreshed
before it can gate anything.
14
15     The first controlled alpha should be a hosted, authenticated product loop: upload a
video, process it, select rallies, compile, and download a reel. The owned-model/commercial
rally-detection claim remains a separate NO-GO until the ship-gate target is ratified and the
WASB checkpoint provenance is resolved.
16
17     ## 1. Problem & goal
18
19     The previous launch blockers were spread across roadmap docs and mixed together three
different questions:

```

```

20
21 - **Can an invited tester use the product without SSH?**
22 - **Can the 15-video golden corpus now support honest quality gates?**
23 - **Can we claim or ship an owned/commercial rally detector?**
24
25 This doc separates those into executable workstreams. The goal is to reach an invite-
only alpha where the product is usable end-to-end, the quality claims are honest and
reproducible, and the remaining model/legal gates are explicit rather than accidentally implied
by launch activity.
26
27 ## 2. Decisions locked
28
29 | # | Decision | Source / date | Implication |
30 |---|---|---|---|
31 | D1 | **Product alpha and owned-model launch are separate decisions.** |
`NEXT_STEPS.md` NO-GO banner, 2026-06-30; verified 2026-07-03 | We can launch a controlled
hosted product loop without claiming the owned model is commercially cleared. |
32 | D2 | **First alpha must be page-driven, not SSH/CLI-driven.** | `NEXT_STEPS.md` P0 UI-
driven/shareable item; verified 2026-07-03 | Hosted service, edge auth, upload, job polling,
compile, and download are launch-critical. |
33 | D3 | **The 15-video corpus is now a measurement asset, not a free launch pass.** |
`data_pipeline/GOLDEN_VIDEOS.md`; local eval runs, 2026-07-03 | Rebase/portability, feature
completeness, baseline refresh, and gate ratification come before quality claims. |
34 | D4 | **Do not promote served Gate A by default from the current n=15 evidence.** |
`backend.eval.serve_contrast` run, 2026-07-03 | Proposal-side lift does not survive served-
windowing safety; keep it as an opt-in/analysis lever. |
35 | D5 | **P10/P11 corpus promotion and train/re-eval are alpha-plus unless owner reopens
D13.** | sibling `khelsutra/docs/STUDIO_MEDIA_LIFECYCLE.md`; verified 2026-07-03 | First alpha
can use manual/offline corpus promotion; automated flywheel follows after serving and gates are
stable. |
36
37 ## 3. Verified snapshot, 2026-07-03
38
39 All five Git checkouts were fast-forwarded before analysis:
40
41 | Repo | Branch | Synced HEAD |
42 |---|---|---|
43 | `badminton-highlight-indexer` | `master` | `e4914d1` |
44 | `khelsutra` | `master` | `b0f70a5` |
45 | `rally-annotator` | `main` | `ef9e088` |
46 | `sports-data-collector` | `master` | `bfc014f` |
47 | `sports-obsreport` | `main` | `a0c7252` |
48
49 Corpus/eval facts:
50
51 - `[verified]` `docs/data_pipeline/GOLDEN_VIDEOS.md` says
`output/human_lovo_manifest.json` is the ingested golden corpus at `n = 15`.
52 - `[analysis]` The local manifest had stale absolute paths. A temporary rebased manifest
found all 15 trajectories and labels in this workspace.
53 - `[analysis]` There are 15 `output/*_golden_features.json` files, but only 6 carry full
shuttle/person/audio candidate features. The newer 9 are windowing-only schema and are skipped
by fusion/person eval.
54 - `[analysis]` `python -m backend.eval.fusion_compare --features-dir output` loaded only
6 videos and measured person-fusion `Delta F1 = -0.021`; not promotable.
55 - `[analysis]` `python -m backend.eval.serve_contrast --manifest <rebased-n15>` measured
proposal Gate A mean `Delta F1 +0.0287`, but served mean `Delta F1 -0.0328`; not served-safe.
56 - `[analysis]` `python -m training.gen0.harness --manifest <rebased-n15> --floor
eval_baselines/heuristic_lovo_n6.json` produced learned LOVO mean F1 `0.551`, `ship=False`;
blockers remain stale n=6 floor/sign test, uncalibrated ship target, and WASB C3 provenance.
57 - `[analysis]` `python scripts/nightly_regression.py --manifest <rebased-n15>
--features-dir output --no-report` reported default/milestone tolF1 `0.3636` vs baseline-off
tolF1 `0.1382`, but marked the frozen baseline stale due to config/corpus changes.
58 - `[analysis]` `python -m backend.eval.ablation --features-dir output --granularity
group` currently fails at import due to a circular import between `backend/eval/ablation.py` and
`backend/eval/ablation_pdf.py`.

```

```

59
60     ## 4. Launch strategy
61
62     ### 4.1 Alpha lane: make the product usable
63
64     The alpha lane is successful when an invited tester can open the app, authenticate,
upload a video, watch the job progress, pick rallies, compile a reel, and download it. This lane
should avoid model-quality promises beyond the measured state and should default to the most
reliable serving path.
65
66     ### 4.2 Quality lane: make the 15-video corpus gate-ready
67
68     The quality lane is successful when the 15-video corpus is portable, reproducible,
feature-complete, and backed by refreshed baselines. Only then should R2 tolF1,
Gen-0/TemporalMaxer, or other model/promotion decisions be ratified.
69
70     ### 4.3 Flywheel lane: make tester feedback useful
71
72     The flywheel lane is successful when user corrections can flow into the collector/vault
path with consent/provenance and a visible re-eval delta. This is valuable for alpha learning,
but it should not block the first controlled alpha if manual promotion can bridge the gap.
73
74     ## 5. Risks and guardrails
75
76     | Risk | Why it matters | Guardrail |
77     |---|---|---|
78     | Launch activity implies model launch approval | The engine docs explicitly say rally-
detection product launch is NO-GO until ship target and WASB C3 clear | Keep all alpha copy
scoped to invite-only processing, not owned-model claims |
79     | Stale local manifests create false confidence | Manifest-driven tools skipped or
failed until paths were rebased | Make the corpus manifest portable and add a path/fixture smoke
check |
80     | 15-video corpus is only partially feature-complete | Fusion/person/audio eval is still
n=6 today | Regenerate or upconvert full-fusion features for all 15 before fusion decisions |
81     | Baseline refresh accidentally rubber-stamps a regression | Nightly runner correctly
reported stale baselines | Refresh baselines only after reviewer sign-off and record the
command/output |
82     | P10/P11 expands scope before hosted alpha is usable | Corpus flywheel is valuable but
can consume the alpha launch window | Keep P10/P11 behind explicit owner reopen unless alpha
lane is already stable |
83     | WASB checkpoint provenance remains unresolved | Blocks commercial sharing of a trained
model regardless of F1 | Treat C3 as a hard model-release gate; do not hide it inside product
launch tasks |
84
85     ## 6. Honest limits
86
87     Completing this doc's alpha lane does **not** prove the owned model is ready, that WASB
weights are commercially clean, that the product can serve unbounded public traffic, that multi-
sport is ready, or that every tester correction automatically improves the model.
88
89     Completing the quality lane does **not** by itself deploy a service. It only makes the
measurement surface honest enough for a launch decision.
90
91     ## 7. Deliverables & progress tracker
92
93     Legend: TODO = not started, IN PROGRESS = actively being worked, DONE =
shipped/verified, GATED = blocked on owner/legal/platform decision. One PR per row unless the
row explicitly names a cross-repo batch.
94
95     | ID | Deliverable | Repo / owner | Depends on | Gated? | Status | PR |
96     |---|---|---|---|---|---|---|
97     | A0 | Create this alpha readiness tracker and index it in docs | `badminton-highlight-
indexer` | - | No | IN PROGRESS | #451 |
98     | A1 | Stand up alpha serving endpoint: Cloud Run/GPU or approved host, health check,
edge auth ON, upload -> process -> jobs -> compile -> download verified | `khelsutra` +

```



```

`badminton-highlight-indexer` | A0 | Owner spend / CF config | TODO | - |
99      | A2 | Make the n=15 golden manifest portable/reproducible; add a path/corpus smoke
check that fails loudly when labels or trajectories are missing | `badminton-highlight-indexer`
| A0 | No | TODO | - |
100     | A3 | Promote the 15-video source-video/MD5 records into the collector/vault path;
reconcile collector's six-video source manifest with the 15-video eval corpus | `sports-data-
collector` + vault | A2 | Owner upload/storage scope | TODO | - |
101     | A4 | Refresh the n=15 nightly regression baseline intentionally and record the exact
command/output; keep stale-baseline warnings until reviewed | `badminton-highlight-indexer` | A2
| Reviewer sign-off | TODO | - |
102     | A5 | Recompute the heuristic n=15 floor and replace the stale `heuristic_lovo_n6`
floor for future Gen-0 comparisons | `badminton-highlight-indexer` | A2 | No | TODO | - |
103     | A6 | Fix `backend.eval.ablation` CLI circular import, then run the R2/tolF1 ablation
on the n=15 corpus | `badminton-highlight-indexer` | A2, A4 | No | TODO | - |
104     | A7 | Regenerate or upconvert full-fusion feature JSONs for the newer 9 videos so
fusion/person/audio analyses run over all 15 | `badminton-highlight-indexer` | A2 | GPU/data
availability | TODO | - |
105     | A8 | Re-run `fusion_compare`, group ablation, and served contrast after A7; explicitly
promote/reject person fusion, Gate A served-default, and any other R-series default flips |
`badminton-highlight-indexer` | A6, A7 | Reviewer sign-off | TODO | - |
106     | A9 | Run Gen-0/TemporalMaxer shadow training against the refreshed n=15 floor; produce
a non-serving artifact bundle and gate verdict | `badminton-highlight-indexer` | A5, A8 | Owner
approval for M0.4 scope | TODO | - |
107     | A10 | Decide and record the alpha ship-gate target: metric, corpus slice, minimum
threshold, significance rule, and what claims are allowed | Owner + `badminton-highlight-
indexer` | A4, A5, A8 | Owner decision | GATED | - |
108     | A11 | Resolve or explicitly defer WASB C3 provenance for alpha messaging; if deferred,
document "no commercial model sharing" copy guardrail | Owner/legal | A10 | Legal/provenance |
GATED | - |
109     | A12 | Reopen or defer P10/P11 corpus flywheel: `PATCH /api/segments`, collector
promotion, `train|reeval` job, and delta-F1 surface | `khealsutra` + `badminton-highlight-
indexer` + collector | A1, A3 | Owner D13 reopen | GATED | - |
110     | A13 | Alpha launch review: runbook, privacy/consent copy, known limits,
rollback/disable steps, and owner GO/NO-GO | Cross-repo | A1, A4, A10, A11 | Owner launch
decision | TODO | - |
111
112     ## 8. Concrete next PR sequence
113
114     1. **Docs PR:** land this tracker.
115     2. **Corpus portability PR:** rebase/portable manifest tooling plus missing-file smoke
test.
116     3. **Hosted alpha PR/batch:** health/auth/upload/process/job/compile/download
verification and runbook updates.
117     4. **Baseline PR:** intentional n=15 nightly baseline refresh and heuristic floor
refresh.
118     5. **Ablation/fusion PRs:** fix ablation import, complete 15-video features, rerun
promotion/rejection decisions.
119     6. **Launch decision PR/docs update:** record ship-gate target, C3 status, alpha copy
limits, and launch checklist.
120
121     ## 9. Open questions
122
123     1. **Owner:** Is the first alpha allowed to use manual/offline corpus promotion, or
should P10/P11 be reopened before any tester wave?
124     2. **Owner/platform:** Which hosted compute target is the next stand-up: Cloud Run GPU,
GCE L4 VM, or another approved GPU host?
125     3. **Owner/legal:** What is the acceptable alpha wording while WASB C3 remains
unresolved?
126     4. **Engineering:** Should the portable corpus manifest live as a generated artifact, a
checked-in redacted manifest, or both?
127     5. **Engineering:** Should n=15 feature regeneration target full fusion schema for every
clip, or should the eval loaders support mixed schemas with explicit capability tags?
128
129     ## 10. Definition of done
130

```

```

131 - [ ] Invited tester can complete the hosted record/upload -> reel download CUJ without
SSH.
132 - [ ] Alpha service is gated behind the chosen auth posture and has a documented
disable/rollback path.
133 - [ ] n=15 corpus is portable enough that manifest-driven quality commands do not rely
on stale local absolute paths.
134 - [ ] n=15 nightly baseline and heuristic floor are refreshed with recorded commands.
135 - [ ] R2/tolF1 gate is promoted or rejected with a recorded ablation result.
136 - [ ] Fusion/person/Gate-A served-default decisions are explicitly promoted or rejected
on the 15-video corpus.
137 - [ ] Ship-gate target and allowed alpha claims are recorded.
138 - [ ] WASB C3 status is recorded and reflected in alpha messaging.
139 - [ ] P10/P11 are either reopened with a committed sequence or explicitly deferred for
alpha-plus.
140 - [ ] Owner records GO/NO-GO for the first alpha wave.
141
142 ## 11. References
143
144 Internal docs:
145
146 - `docs/NEXT_STEPS.md`
147 - `docs/ROADMAP.md`
148 - `docs/data_pipeline/GOLDEN_VIDEOS.md`
149 - `docs/R2_EVAL_METRIC_DESIGN.md`
150 - `docs/RALLY_DETECTION_FIXES_PLAN.md`
151 - `docs/COMPUTE_DECOUPLED_SERVING/README.md`
152 - sibling `khehsutra/NEXT_STEPS.md`
153 - sibling `khehsutra/docs/STUDIO_MEDIA_LIFECYCLE.md`
154 - sibling `khehsutra/docs/PER_RALLY_SELECTION.md`
155 - sibling `sports-data-collector/data/golden_source_videos.json`
156
157 Local verification commands used for the 2026-07-03 snapshot:
158
159 ```powershell
160 python -m backend.eval.fusion_compare --features-dir output
161 python -m backend.eval.serve_contrast --manifest <rebased-n15-manifest>
162 python -m training.gen0.harness --manifest <rebased-n15-manifest> --floor
eval_baselines\heuristic_lovo_n6.json --version analysis-n15 --out <temp-out>
163 python scripts\nightly_regression.py --manifest <rebased-n15-manifest> --features-dir
output --no-report
164 python -m backend.eval.ablation --features-dir output --granularity group
165 ```
166
167 ---
168
169 ### Changelog
170
171 - 2026-07-03 - Created from a synced five-repo review of `NEXT_STEPS`, `ROADMAP`,
golden-corpus docs, local eval artifacts, and the current 15-video corpus state.
172

```

10. assistant (2026-07-03T16:40:37.940Z)

Now let me check the archives/decisions directory to understand the ADR convention:

11. user (2026-07-03T16:40:38.626Z)

```

1 # Architecture Decision Records
2
3 ## ADR-001: Precise shuttle tracking via self-hosted TrackNetV4/WASB on WSL2
4
5 **Date:** 2026-05-25
6 **Status:** Accepted
7
8 ### Context

```

```

9
10 The pipeline uses an AI multimodal provider (Gemini) for *semantic* rally segmentation.
11 We need **precise frame-level detection of the moving shuttlecock** (and, generically,
12 balls for other sports) to produce higher-quality candidate rallies and a fine-tuning
13 signal.
14
15 We evaluated whether a commercially available model/API could do this. Findings:
16
17 - **Multimodal LLM APIs (Gemini, Claude, GPT-class) are the wrong tool for precise
18 localization.** They reason well about *what* is happening (rally vs. dead time) but
19 cannot reliably track a tiny, fast (~up to 400 km/h), motion-blurred shuttle frame by
20 frame.
21 - The precision leaders are **purpose-built trajectory models**, which are open source
and
22 self-hosted, not big-vendor APIs:
23 - **TrackNetV4** (TensorFlow) ? badminton-specific, heatmap + motion-attention,
24 highest reported shuttle accuracy; ships `src/predict.py` producing per-frame (X, Y)
25 coordinate CSV. https://github.com/AR4152/TrackNetV4
26 - **WASB** (PyTorch) ? multi-sport strong baseline (badminton/tennis/volleyball/?),
27 fits our `SportProfile` abstraction; needs a custom-video inference wrapper.
28 https://github.com/nttcom/WASB-SBDT
29 - True turnkey commercial APIs (Roboflow hosted models) exist but cap below a tuned
30 TrackNet; enterprise systems (Hawk-Eye, SkillCorner, TRACAB) are not API-accessible.
31
32 ### Decision
33
34 1. Adopt **self-hosted TrackNetV4** as the primary precise shuttle detector, with
**WASB**
35 as the multi-sport fallback.
36 2. Run these models under **WSL2 (Ubuntu)**, **not** native Ubuntu. The entire indexer
37 already runs on Windows; WSL2 keeps a single dev environment and quarantines the
38 TensorFlow/PyTorch dependencies. Native Ubuntu would require dual-boot and migrating
39 the whole project for marginal gain.
40 3. Integrate via a **subprocess adapter**: a detector module that shells out to the WSL
41 `predict.py`, reads the coordinate CSV, and converts trajectory ? candidate rallies
42 feeding the existing stitching flow. This keeps the TF dependency fully isolated from
43 the main Python environment.
44
45 ### Consequences
46
47 - **WSL I/O caveat:** `/mnt/c` access is slow. Stage match videos inside the WSL
48 filesystem (e.g. `~/clips/`) before inference rather than reading large mp4s across
the
49 Windows?WSL boundary.
50 - TrackNetV4 (TensorFlow) and WASB (PyTorch) live in **separate conda envs**.
51 - Next step after integration: manual video annotation ? fine-tune the detector.
52
53 ---
54
55 ## ADR-002: Use the sports-data-collector's rally annotations as a golden eval +
temporal-segmenter training set
56
57 **Date:** 2026-05-26
58 **Status:** Accepted
59
60 ### Context
61
62 A sibling project, **sports-data-collector**, curates sports videos and stores per-rally
63 `[start, end)` timestamps (plus stroke counts, scores) in an SQLite DB, with a
64 commercial-license flag per source. We asked whether this data could improve the
65 indexer's rally detection, and for which task.
66
67 Key findings:
68
69 - The annotations are **temporal-only** (when rallies start/end). They do **not**

```

```

contain
70     per-frame shuttle `(x, y)` coordinates, so they **cannot train the shuttle detector**
71     (TrackNetV4/WASB ? see ADR-001), which needs coordinate labels.
72     - They are an excellent **golden evaluation set** and a **training signal for the
73     **temporal rally segmenter** (active vs. dead play over time) ? the stage that is
today
74     a hand-tuned heuristic, not a learned model.
75     - The collector's five sports tables (`badminton_rallies`, `tennis_points`,
76     `cricket_deliveries`, `pickleball_rallies`, `tabletennis_rallies`) all share a
77     `(video_id, ordinal, start_time, end_time)` shape, and the indexer already has
78     `sport_registry` / `segmenter_registry`. So a *generic* `(start, end)` bridge serves
79     every sport.
80     - The two repos key videos differently: the collector by a human `video_id`, the indexer
81     by **file MD5**. A re-encode/re-download changes the bytes, hence the MD5.
82
83     ### Decision
84
85     1. Treat the collector's rally timestamps as (a) the **golden eval set** for rally
86     *segmentation* and (b) the **labels** for training a temporal segmenter ? explicitly
87     **not** as shuttle-detector training data.
88     2. Bridge the repos with a sport-agnostic **export** (`curate export` in the collector):
89     per-video `start,end` CSVs (the indexer's eval format) + a `manifest.json` that pairs
90     each CSV with its video file. Gate on `is_commercial` so only commercially usable
data
91     flows through.
92     3. Honour the **MD5 keying** contract: evaluation/processing must run on the exact video
93     file recorded in the manifest; acquisition (ADR-003) therefore precedes processing.
94
95     ### Consequences
96
97     - The collector becomes the *producer* of labels; the indexer the *consumer* (processing
+
98     scoring + training). Neither takes a hard dependency on the other beyond the manifest.
99     - Scoring is free and deterministic (reads the DB; no API). See `docs/EVALUATION.md`.
100    - Only invest in frame-level coordinate labels (for the detector) **if** evaluation
shows

```

12. assistant (2026-07-03T16:40:41.120Z)

Let me check for any existing conventions about decision recording in
DECISION_SNAPSHOT_REGRESSION.md or similar docs:

13. user (2026-07-03T16:40:41.689Z)

```

1      <!--
2      PROJECT DOC ? tracked design plan for a deterministic "quality-neutral proof".
3      Follows docs/PROJECT_DOC_TEMPLATE.md. $7 progress tracker is the source of truth.
4      -->
5
6      # Decision-Snapshot Regression ? a deterministic quality-neutral proof
7
8      > **Status:** `DRAFT` ? **$2 decisions LOCKED (2026-06-25); ready to start $7 P1** .
**Owner:** Avi . **Created:** 2026-06-25 . **Last updated:** 2026-06-25
9      > **Lifecycle:** `DRAFT` ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE)
10     > **Tracking anchors:** $7 progress tracker is the source of truth; indexed in
`docs/README.md`; pointer in memory `current-state` once work starts.
11     > **Relation to existing docs:** **extends** `GOLDEN_REGRESSION_FIXTURES.md` (post-
detector fixtures) and the existing characterization snapshot in
`tests/test_golden_fixtures.py`; **peer of** the M4 quality-floor track and the GPU
`scripts/nightly_reelcount.py` e2e nightly.
12     > **Honesty note:** claims tagged `[verified]` (checked vs code) / `[design]`
(proposed).
13
14     ---

```

```

15
16     ## 0. TL;DR
17     A deterministic regression that proves a change is quality-neutral (behavior-
preserving) by freezing and comparing the pipeline's decision output ? the rally segment
list + the stitch/concat plan ? run off a cached trajectory (GPU-free). It extends
the existing characterization snapshot (`test_replay_matches_committed_expected`, today
segmentation-only) down to the stitch plan, runs offline in CI, and gives the behavioral
proof that lets a *genuine* refactor (not just `ruff format`) be confidently called quality-
neutral. Most important caveat: never hash the encoded reel `mp4` ? ffmpeg/x264/NVENC/mux-
timestamps make those bytes non-deterministic, so a reel-MD5 would false-alarm on every ffmpeg
bump.
18
19     ## 1. Problem & goal
20     "Presumably quality-neutral" changes (refactors) need a cheap, deterministic proof that
they actually preserve behavior ? so they can be merged (and auto-signed-off) with confidence.
Today we have two weaker tools: AST-identity [verified] (used to bless `ruff format` in
#379) only proves *formatting* is unchanged ? it cannot bless a real refactor (function
extraction, rename, loop reflow) that changes the AST but should preserve behavior; and the full
real-video F1 / reel-count e2e needs GPU + the gitignored corpus
(`scripts/nightly_reelcount.py`) [verified] so it can't gate a PR. The decision-snapshot fills
the gap: a deterministic, offline, behavior-level proof.
21
22     What "good" looks like: a refactor PR runs the snapshot in CI; byte-identical
decision output ? provably quality-neutral ? eligible for auto-sign-off; a behavior-changing
edit fails with a readable diff.
23
24     Read first: `tests/test_golden_fixtures.py::test_replay_matches_committed_expected`
[verified] (the existing snapshot ? its own message: "behaviour changed vs frozen snapshot
(characterization, not correctness)"), `backend/eval/golden_fixtures.py`,
`backend/pipeline/stitcher/compiler.py` (the concat plan), audit doc B1 (`motion`'s random
fields), `scripts/nightly_reelcount.py` (the distinct GPU e2e).
25
26     ## 2. Decisions *(D1?D11 **LOCKED** 2026-06-25)*
27
28     | # | Decision (recommendation) | Rationale | Status |
29     |---|-----|-----|-----|
30     | D1 | Snapshot the DECISION output, never the encoded reel. Freeze the rally
segment list + the stitch/concat plan; never MD5 the reel `mp4`. | Encoded mp4 bytes are non-
deterministic (ffmpeg version string, x264 threading, NVENC, mux timestamps) ? a reel-MD5 false-
alarms on every ffmpeg bump/machine. The repo's own `nightly_reelcount.py` already asserts
count+metrics, not bytes [verified]. | LOCKED (2026-06-25) |
31     | D2 | Input = a cached trajectory (detector output), not raw video. Run from a
pinned trajectory CSV (the `fx_r0X` fixtures). | The WASB detector (video?trajectory) is the
GPU/WSL + CV-float-non-deterministic upstream; caching its output makes the proxy deterministic
+ CI-runnable [verified ? the fixtures are post-detector]. | LOCKED (2026-06-25) |
32     | D3 | Coverage = segmentation decision + stitch/concat plan; STOP before the
encode. Extend the existing segmentation snapshot down to the concat manifest (clips, cut
points, order, padding). | Segmentation is already snapshotted [verified]; the stitch plan is
the deterministic missing layer. The encode is non-deterministic (D1). | LOCKED (2026-06-25)
|
33     | D4 | Compare = canonical-JSON snapshot with a readable diff (an MD5 of the
canonical JSON may be a fast-path summary, but the field-level diff is the artifact). | A
refactor that trips it needs a debuggable diff, not "hash changed" ? mirrors
`test_replay_matches_committed_expected`'s mismatch report [verified]. | LOCKED (2026-06-25)
|
34     | D5 | Snapshot only deterministic segmenters/fields; exclude the random
fabricated fields. Use a deterministic segmenter (CONTROL/fusion on trajectory) and a
`{start,end,shuttle}`-style schema that excludes `motion`'s random
server/receiver/shot_count/avg_speed`. | `motion` fabricates random fields (audit B1)
[verified] ? non-snapshot-able as-is. B1's fix (drop the fabrication) is the enabler that
would make `motion` snapshot-able. | LOCKED (2026-06-25) |
35     | D6 | The reel mp4 gets a FUNCTIONAL smoke, not a byte hash (exists, ~duration
within tol, N segments, sane streams). | Corollary of D1 ? verify the encode is *sane*, not
*identical*. | LOCKED (2026-06-25) |
36     | D7 | Runs in per-PR CI, offline, CONTROL-only ? a refactor gate, not a quality

```

```

floor.** | The point is catching a behavior change in a refactor *before merge*; it must be
fast, GPU-free, on the pinned CONTROL config. | **LOCKED** (2026-06-25) |
37 | **D8** | **Intentional changes re-freeze WITH a documented rationale (Launch
Report).** A `--update-snapshot` flag exists, but a mismatch on an *intended* change must be re-
frozen with the Launch-Report rationale recorded. | Without the discipline a snapshot becomes a
rubber stamp; ties to the launch-report convention. | **LOCKED** (2026-06-25) |
38 | **D9** | **Scope boundary: this is the quality-NEUTRAL proxy ? distinct from the M4
quality-FLOOR (per-slug F1), Tier-2 real-video e2e (GPU), and `nightly_reelcount.py`.** Keep
them separate + complementary. | Conflating "behavior unchanged" with "quality is good enough"
muddies both. | **LOCKED** (2026-06-25) |
39 | **D10** | **Deterministic-segmenters-only, ENFORCED by a determinism guard.** Cover
only deterministic segmenters; to *ensure* no non-deterministic one slips in (? a flaky
snapshot), the harness runs each fixture **twice and asserts byte-identical decision output**
before freezing ? any segmenter whose two runs differ (hidden `random`, set-ordering, float/time
non-determinism) is **rejected from snapshot scope** and flagged. Mechanical, not a hand-
maintained allowlist. | Owner, 2026-06-25. `motion` is non-deterministic until **B1 / #385**
lands; then it auto-becomes eligible (no allowlist edit). | **LOCKED** (2026-06-25) |
40 | **D11** | **Auto-sign-off integration is STAGED: ADVISORY first ? GATING later.** Ship
the snapshot as an advisory/informational check; graduate it to a **gating** check (gates a non-
AST-identical refactor's auto-LGTM) only after it's proven stable on real refactors. | Owner,
2026-06-25. Resolves S8 Q1 ? the strategic payoff, de-risked by starting advisory. | **LOCKED**
(2026-06-25) |
41
42 ## 4. Design / architecture
43 **Flow (all offline, deterministic):**
44 `pinned trajectory CSV ? [segmentation: CONTROL segmenter] ? segment list ?
[deterministic stitch-PLAN builder: clips + cut points + padding + order, NO ffmpeg] ? canonical
JSON {schema_version, segments[], stitch_plan{}} ? MD5 / field-diff vs frozen golden`.
45 Mismatch ? behavior changed (decide if intended; re-freeze via D8).
46
47 **Reuse map** [verified unless noted]:
48 - `tests/test_golden_fixtures.py::test_replay_matches_committed_expected` +
`backend/eval/golden_fixtures.py` ? the existing characterization snapshot at the
**segmentation** layer; extend its `expected` schema + replay to also carry the stitch plan.
**(reuse + extend)**
49 - `eval_baselines/fixtures/` (`fx_r0X` + `manifest.json`) ? the pinned post-detector
trajectories. **(reuse)**
50 - `backend/pipeline/stitcher/compiler.py` ? builds the ffmpeg concat today; **[design]**
factor out a pure `build_stitch_plan(segments, cfg) -> dict` (the *decision*) from the *encode*,
so the plan is hashable without running ffmpeg. **(new seam)**
51 - `eval_baselines/` snapshot store + the cross-OS enforcement already used by the golden
suite. **(reuse)**
52
53 ## 6. Honest limits ? what this does NOT do
54 - Does **not** prove *quality is good* ? only that *behavior is unchanged* vs the frozen
golden (that's M4 / F1-floor's job).
55 - Does **not** cover the **encode** (pixels/bytes) ? only the decision + stitch plan;
the reel gets a functional smoke only (D6).
56 - Does **not** run the **real detector** (GPU/WSL) or real videos ? it runs from cached
trajectories (Tier-2 / `nightly_reelcount.py` cover that, GPU-gated).
57 - A **green snapshot does not bless a behavior change** ? it only certifies *no* change;
an intended change must be re-frozen with a rationale (D8).
58 - Only as good as the **fixtures' coverage** ? pipeline paths not exercised by `fx_r0X`
aren't protected.
59
60 ## 7. Deliverables & progress tracker ? **source of truth**
61 Legend: ? Todo . ? In progress . ? Done . ? Blocked/gated. One small PR per row,
branched from `origin/master`.
62
63 | ID | Deliverable | Depends on | Gated? | Status | PR |
64 |----|-----|-----|-----|-----|----|
65 | P0 | This doc ? decisions in S2 locked | ? | Owner | ? | (this PR) |
66 | P1 | Canonical decision-snapshot schema + deterministic serializer (segments + stitch
plan; excludes non-det fields) | P0 | No | ? | ? |
67 | P2 | Factor `build_stitch_plan()` out of `compiler.py` (pure, ffmpeg-free) + unit test

```

```

| P0 | No | ? | ? |
68 | P3 | Extend `golden_fixtures` / `test_golden_fixtures` to capture + compare the stitch
plan **+ the D10 determinism guard (run each fixture twice ? assert byte-identical; reject non-
deterministic segmenters)** | P1, P2 | No | ? | ? |
69 | P4 | Freeze CONTROL snapshots for the existing `fx_r0X` fixtures | P3 | No | ? | ? |
70 | P5 | `--update-snapshot` flow + Launch-Report hook (D8) | P3 | No | ? | ? |
71 | P6 | Wire the snapshot as a per-PR CI characterization gate (CONTROL-only) | P4 | No |
? | ? |
72 | P7 | **Surface the snapshot on refactor PRs ? ADVISORY** (informational; does not
gate) ? per D11 | P6 | No | ? | ? |
73 | P8 | **(gated)** Graduate to GATING auto-sign-off: a snapshot-match marks a non-AST-
identical refactor quality-neutral ? LGTM-eligible (only after P7 proven stable) ? per D11 | P7
| **Owner** | ? | ? |
74
75 ## 8. Open questions ? owner
76 1. ? **RESOLVED (D11):** advisory first ? graduate to gating. *(The strategic payoff,
de-risked by staging.)*
77 2. ? **RESOLVED (D10):** deterministic-segmenters-only, **enforced by the double-run
determinism guard** ; **B1 is an enabler, not a blocker** ? and it's being fixed in **385**,
which makes `motion` auto-eligible.
78 3. **(still open) Fixture scope:** reuse the existing `fx_r0X` (which stop at
segmentation) by adding a stitch-plan `expected`, or add a dedicated fixture set?
79 4. ? **DONE: D1?D11 LOCKED 2026-06-25.**
80

```

14. assistant (2026-07-03T16:40:43.785Z)

Now let me check what changes were made to DOC_STATUS.md and NEXT_STEPS.md in this PR:

15. user (2026-07-03T16:40:45.201Z)

```

diff --git a/docs/DOC_STATUS.md b/docs/DOC_STATUS.md
index 34411fd..2541672 100644
--- a/docs/DOC_STATUS.md
+++ b/docs/DOC_STATUS.md
@@ -30,6 +30,7 @@ Before a foundational/tracked doc moves to `docs/archives/`:
### 3a. Tracked projects ? lifecycle docs (archive when DONE; each doc's $7 is the completeness
source-of-truth)
| Doc | Status | Archive-ready? | Next action |
| ---|---|---|---|
+| `ALPHA_LAUNCH_READINESS.md` | **DRAFT** ? cross-repo alpha readiness tracker created
2026-07-03 after the golden corpus reached n=15; separates hosted product-alpha work from owned-
model/commercial rally-detection gates | No | Execute $7 in order: corpus portability + n=15
baselines, hosted alpha serving, ablation/fusion reruns, then launch GO/NO-GO |
| `COMPUTE_DECOUPLED_SERVING/` | **IN PROGRESS** ? M0 approved + **M1?M9 LARGELY SHIPPED**
(#279?#290; M5 #286, M6 #287, M8 #288, M9-auth #290 all merged, default-OFF); Gen-0 weights
pushed (`weights/0.1.0-l4/`) | No | M11 flywheel + M12 gdrive-api (owner residual:
counsel/prices/OAuth/real-run within the D17 cap) |
| `RALLY_QUALITY_RESEARCH.md` | Foundation, **tracked & ACTIVE** ($7) | No | Drives the
R-series; complete when $7.1 single quality number is attributed |
| `R2_EVAL_METRIC_DESIGN.md` | IMPLEMENTED (augmenting) | No (gates R4 + dual-eval; ablation
pending) | Promote-or-reject toF1 as the gate (ablation) ? then archive with the R-cluster |
diff --git a/docs/NEXT_STEPS.md b/docs/NEXT_STEPS.md
index f186ff9..a9f35a8 100644
--- a/docs/NEXT_STEPS.md
+++ b/docs/NEXT_STEPS.md
@@ -9,7 +9,7 @@
## ? Prioritized pending ? all tracks
**P0 ? highest leverage / unblocks the most**
-0. ** ? [owner-priority, 2026-06-21] MAKE IT UI-DRIVEN + ALPHA-SHAREABLE ? "a few recordings
lying around ? highlights, from a page, no SSH".** Today the cloud flow is **SSH/CLI-only**
(`docs/CLOUD_GPU_DRYRUN_GUIDE.md`); **M12 does NOT fix it.** Needs (a separate khelsutra `web/`
SPA + deploy track, NOT an engine seam): *(i)** deploy the engine as a reachable **service**
(M7 ? Cloud Run service, ? owner GCP spend for a *persistent* endpoint), *(ii)** the **SPA

```

ingest/progress screen** (khelsutra `web/` B-P2: enter/upload video ? `/api/process` ? poll `/api/jobs` ? reel; ZERO engine code), ** (iii)** flip **M9 edge-auth ON** (merged #290, default-OFF) to gate alpha tenants, ** (iv)** a **compute PICKER in the UI** ("your machine vs our cloud", D13/D15). The engine seams make it *possible*; the SPA makes it *usable*. Full record: [[ui-driven-cloud-serving-gap]].

+0. **? [owner-priority, 2026-06-21] MAKE IT UI-DRIVEN + ALPHA-SHAREABLE ? "a few recordings lying around ? highlights, from a page, no SSH".** Today the cloud flow is **SSH/CLI-only** (`docs/CLOUD\_GPU\_DRYRUN\_GUIDE.md`); **M12 does NOT fix it.** Needs (a separate khelsutra `web/` SPA + deploy track, NOT an engine seam): ** (i)** deploy the engine as a reachable **service** (M7 ? Cloud Run service, ? owner GCP spend for a *persistent* endpoint), ** (ii)** the **SPA ingest/progress screen** (khelsutra `web/` B-P2: enter/upload video ? `/api/process` ? poll `/api/jobs` ? reel; ZERO engine code), ** (iii)** flip **M9 edge-auth ON** (merged #290, default-OFF) to gate alpha tenants, ** (iv)** a **compute PICKER in the UI** ("your machine vs our cloud", D13/D15). The engine seams make it *possible*; the SPA makes it *usable*. Full record: [[ui-driven-cloud-serving-gap]]. **Concrete execution tracker after the n=15 corpus review:** `docs/ALPHA\_LAUNCH\_READINESS.md` §7.

0b. **? [packaging track, IN PROGRESS ? 11 PRs merged 2026-06-21/22] Downloadable batteries-included app ? `docs/PACKAGING\_PLAN.md`.** The LIGHT-tier engine is bundle-ready: `pip install` ? `indexer --app` serves the **bundled SPA** single-origin, **torch-free**, with app-home state / ffmpeg resolver / DNS-rebinding guard / single-instance lock / safe DB upgrades (P0a?e, P1a/b, P2-launcher/P2a/P2d/P2e). **? Next (fresh focused pushes):** ** (1)** P3 **GPU validation + screencast** (L4, ~\$0.50, DELETE-after, \$100 cap ? `[[gcp-vm-always-delete]]` ; ? native-WASB 0-rally tuned-config blocker), ** (2)** P2b-install (uv/script + truly-local default config), ** (3)** P2c first-run wizard (cross-repo khelsutra SPA from `/api/capabilities`), then P3-cloud (Cloud Run) + P4 (ONNX/train/annotate = north star). Memory `[[packaging-app-plan]]` / `[[next-session-packaging-app]]`.

1. ? **`COMPUTE\_DECOUPLED\_SERVING` M0?M9 LARGELY SHIPPED** (2026-06-21, ADR-014 ratified). M1?M4 (#279/#280/#282/#283), **M5 #286, M6 #287, M8 #288, M9-auth #290** all merged (default-OFF); the **real M7 L4 run** validated the worker (22 rallies, VM deleted, \$0). All \$8 gating Qs locked (D7?D17). **Owner residual** (not blocking Claude's default-OFF code): counsel ToS/DPA (M9-consent live + M11 live), real GCP prices+FX (M8 calibrate), Google OAuth app (M12 real), CF team-domain/AUD + ingress-lock; real cloud verification authorized within a **\$100/?10k cap** (D17).

2. **[owner, ongoing] Grow the golden corpus ? multi-court badminton + ?3 matches/sport (TT first). The single highest-leverage move; the quality track *and* full-corpus training feed on it.

16. assistant (2026-07-03T16:40:48.611Z)

Now let me check if there are any existing cross-repo trackers or similar patterns:

17. user (2026-07-03T16:40:57.562Z)

```
./badminton-highlight-indexer/docs/ALPHA_LAUNCH_READINESS.md
./badminton-highlight-indexer/docs/archives/DEFERRED_HARDENING_PLAN-COMPLETED-2026-06-14.md
./badminton-highlight-indexer/docs/archives/past_projects/AUDIT_REPORT_khelsutra-
guru_2026-07-02-COMPLETED-2026-07-02.md
./badminton-highlight-
indexer/docs/archives/past_projects/CODE_CLEANUP_AUDIT_2026-06-24-SUPERSEDED-2026-07-02.md
./badminton-highlight-indexer/docs/BACKLOG.md
./badminton-highlight-indexer/docs/DOC_STATUS.md
./badminton-highlight-indexer/docs/NEXT_STEPS.md
./badminton-highlight-indexer/docs/PACKAGING_PLAN.md
./badminton-highlight-indexer/docs/POST_166_WORK_ITEMS_MILESTONES_AND_RISK_ASSESSMENT.md
./badminton-highlight-indexer/docs/README.md
```

18. assistant (2026-07-03T16:40:59.628Z)

Let me check the sibling khelsutra repo docs structure:

19. user (2026-07-03T16:41:00.464Z)

```
total 192
drwxr-xr-x 1 avidu 197609    0 Jul  2 21:13 .
```



```
drwxr-xr-x 1 avidu 197609      0 Jul  2 21:13 ..
-rw-r--r-- 1 avidu 197609 17371 Jul  2 08:13 ANNOTATE_AND_TRAIN.md
drwxr-xr-x 1 avidu 197609      0 Jun 23 00:19 archives
-rw-r--r-- 1 avidu 197609 16474 Jul  2 21:13 DRY_RUN_CUJ.md
-rw-r--r-- 1 avidu 197609 37141 Jul  2 21:13 LOCAL_APPLIANCE_ARCHITECTURE.md
-rw-r--r-- 1 avidu 197609 17451 Jun 23 00:19 MULTILINGUAL_EXPANSION.md
-rw-r--r-- 1 avidu 197609 29326 Jun 22 15:27 PER_RALLY_SELECTION.md
-rw-r--r-- 1 avidu 197609 5000 Jul  2 21:13 PORTABILITY.md
-rw-r--r-- 1 avidu 197609 3790 Jul  2 08:13 README.md
-rw-r--r-- 1 avidu 197609 6997 Jun 22 15:27 REQUEST_ACCESS_FORM.md
-rw-r--r-- 1 avidu 197609 29193 Jul  2 08:13 STUDIO_MEDIA_LIFECYCLE.md
```

20. assistant (2026-07-03T16:41:05.378Z)

Let me check NEXT_STEPS.md to see if there's a convention about cross-repo coordination:

21. user (2026-07-03T16:41:06.154Z)

```
font eyeball) CLEARED;** carried forward to `NEXT_STEPS.md`: a delta native-review of the
post-sign-off
string additions, and **B7** (engine backend l10n, deferred).
- **`MULTILINGUAL_EXPANSION.md`** ? `IN PROGRESS`: **Phase A (es/da/id) shipped** ([#67]);
**Thai
(B) / CJK (C) remain owner-gated**. Adds **7 locales** (th/id/da/es/zh-Hans/ja/ko) beyond the
en/kn/hi foundation with en/kn/hi-level diligence; peer of the archived
`FRIENDLY_AND_MULTILINGUAL.md`, reusing its foundation.
```

22. assistant (2026-07-03T16:41:08.292Z)

Let me check if there's a NEXT_STEPS.md in khelsutra:

23. user (2026-07-03T16:41:09.086Z)

```
-rw-r--r-- 1 avidu 197609 9773 Jul  2 21:13 NEXT_STEPS.md
```

24. user (2026-07-03T16:41:11.088Z)

```
1      # Next steps (live) ? KhelSutra brand / site
2
3      Prioritized "what to do next" for the **`khelsutra` product repo**. Refresh at each
checkpoint.
4
5      **Where we are (updated 2026-07-02):** `khelsutra.guru` is **LIVE and fully shipped** ?
landing page + `/capture`
6      checklist + OG image + a **segmented Request-Access form** (academy/individual + GPU +
storage + free
7      text). On **Cloudflare** (Worker `hidden-king-khelsutra` + D1 `khelsutra-access`): the
form persists to
8      D1 **and emails a summary to the owner on every registration** (Email Routing send
binding ? verified
9      live). **Cloudflare is the single canonical host** ? the DO droplet mirror
(`168.144.126.176`) was
10     **decommissioned + deleted 2026-06-25** (#117; data archived to
11     `gs://khelsutra-rally-corpus/archive/droplet-claude-do-rally-test-20260625/`). The
dependency-free
12     **Node + SQLite** host (`site/server.js`) remains VM-portable (live-proven on the
droplet before its
13     retirement) but nothing is deployed on it today; alpha serving is **paused pending the
Cloud Run move**.
14     **Git auto-deploy is connected** (Workers Builds, root `site/`): merge to `master` ?
deploys; **CI gates
15     every PR** ? three jobs: the full site `npm test` (Worker + VM host + assets + UX), the
`web/` SPA build
16     and coverage floor, and a real-browser layout check. Shipped across ~90 merged PRs (the
```

```

form, 6-language
17     localization, the per-rally MVP, VM portability). Docs: `docs/REQUEST_ACCESS_FORM.md`,
18     `docs/PORTABILITY.md`.
19
20     ## Next product threads (owner ideas, 2026-06-20)
21     - **? DONE ? Friendly & Multilingual (ARCHIVED 2026-06-22; human gates cleared):** the
product is
22         non-tech-friendly and **fully localized to all 6 languages** (en/kn/hi/es/da/id)
across both surfaces
23         (`web/` SPA + `site/` marketing) ? including the home page (hero, all sections, the
`#join` access form,
24         footer), the **Privacy Policy** page, **per-page`<title>`s**, and **es/da/id academy
flyers** (with
25         localized SPA screenshots). The responsive header is fixed and a **headless-Chromium
CI suite** guards
26         layout, cross-locale render (live DOM == catalog), and **accessibility (axe-core WCAG
2.1 AA)**. Full
27         record: **`docs/archives/FRIENDLY_AND_MULTILINGUAL.md`**. Post-archive PRs: #86 (own-
model) · #87 (es/da/id
28         site) · #105 (form+footer + render guard) · #106 (Privacy page) · #107 (a11y fixes +
axe guard + page
29         titles) · #109 (flyers). **Owner gates CLEARED:** native-speaker review of the AI-
drafted catalogs ? +
30         real-device font-glyph eyeball ?. **Remaining (small/owner):**
31         - **Delta native-review** of the strings ADDED after the owner's sign-off
(#105/#106/#107/#109 ? the
32         form/footer/privacy/title/flyer copy is AI-drafted): a small top-up review before a
wide public launch.
33         - **B7 (deferred, not a v1 blocker):** engine backend customer-facing l10n (report
`customer_` +
34         rejection reasons) ? cross-repo Python (`badminton-highlight-indexer`), owner-gated
**D-BACKEND-DEPTH**.
35         - **Marketing video on the site** ? embed a short walkthrough of the record ? index ?
reel process,
36         themed for **serious enthusiasts**, on the landing page. ? **A real-UI CUJ screencast
now exists**
37         (Playwright recording of the per-rally SPA driven against real GPU-run data ?
`promo_cuj.mp4/`.gif`
38         in `gs://khehsutra-rally-corpus/promo/`). Decide: use it as the marketing clip (?
shows real court
39         footage + a dev "P5" badge ? clean up before public embed), and/or produce a polished
Grok/Gemini cut.
40         - **Interactive per-rally selection** (big product feature) ? ? **MVP / M0 COMPLETE**
(the full
41         pick ? query ? select ? compile ? download loop on existing engine endpoints). Lives
in the alpha SPA
42         (`web/`). **? Tracked design + P-tracker (source of truth):
`docs/PER_RALLY_SELECTION.md` (IN PROGRESS;
43         owner-approved, D8?D11 locked).** **Done: P0 design (#11), P1 scaffold (#14), P2
grid+hook (#21), P3
44         `parseQuery()` query bar (#28), P4 SelectionFooter + `min_shot_count` warning (#30),
P5 real load +
45         Build?compile?preview?download (#31).** Verified e2e vs a stub engine + a CUJ replay
against the real
46         L4-run rallies + reel. **NEXT (owner-gated, engine repo): M1?M5** ? `GET /api/health`
+ `download_url`
47         (P6), RESTful `?/rallies` alias (P7), multi-video home + auth (P8), per-rally
preview/timeline (P9),
48         corpus correction loop (P10), localized NL (M5). **Optional:** a fresh live
`/api/process` GPU
49         inference once the engine session's M2 training frees the single GPU.
50
51     ## Queued (site polish)
52     - **Always-Use-HTTPS + HSTS** (Cloudflare SSL/TLS) ? closes the plain-HTTP flag.
53     - **Cloudflare Web Analytics (owner-gated; FREE)** ? code side DONE (the dead

```

placeholder beacon was

```

54     removed from `site/public/index.html`; analytics is now via **Automatic Setup**, no
in-source token).
55     **Owner step (needs CF login):** the site is a Cloudflare **Worker** (not Pages) on
the proxied
56     `khelsutra.guru` zone ? enable via **Analytics & Logs ? Web Analytics ? Add a site ?
`khelsutra.guru`**
57     (Automatic Setup is default-on for proxied hostnames; edge-injects the beacon on every
page). Then
58     filter the view to `hostname = khelsutra.guru` to exclude `*.pages.dev`/preview noise,
and decide
59     separately whether `app.khelsutra.guru` (the SPA) needs its own enablement (history-
router only).
60     **Verify (two gotchas):** the Worker's asset responses must NOT send `Cache-Control:
?, no-transform`
61     (Cloudflare can't rewrite that ? no injection), and after a deploy the rendered HTML
should contain the
62     edge-injected beacon. **DONE WHEN:** the Web Analytics dashboard shows incrementing
page views for
63     `khelsutra.guru`.
64     - **Any future VM ? production** *(re-scoped 2026-07-02 ? was "Droplet ? production";
the droplet was
65     deleted 2026-06-25, #117)* ? only if a VM standby is ever wanted again:
`site/scripts/provision-vm.sh`
66     is provider-agnostic; add TLS + a subdomain (`vm.khelsutra.guru` via Caddy) and **VM-
path SMTP** so the
67     Node host also emails (currently `noopNotify`). Not urgent ? Cloudflare is the single
canonical host.
68     - **Revoke the droplet SSH key (`~/ssh/khelsutra_droplet`) ? still actionable, do it
now (owner):**
69     the droplet is gone (deleted 2026-06-25) but the local keypair lingers; revoke/delete
it rather than
70     waiting on any VM work.
71     - **DMARC** TXT for deliverability; delete any orphan `khelsutra-site` Worker; clean CF
D1 test rows
72     (`deploy-test@`, `email-test@`).
73
74     ## Bigger arcs (design-first)
75     - **? Studio media lifecycle (ingest ? store ? annotate ? train)** ? **Tracked project:
`docs/STUDIO_MEDIA_LIFECYCLE.md`
76     (IN PROGRESS ? storage half COMPLETE 2026-07-01). ** The seamless flow: add a video ?
store it in a user-chosen medium
77     (S3/GCS/Drive/local) ? free-space checked ? annotate in-UI ? contribute to the
segmenter. **SHIPPED 2026-07-01:**
78     free-space guards (vault + engine, P1/P2), the `/api/capabilities` storage block (P3),
`POST /api/assets`
79     store-into-medium (P4), the SPA medium picker (P5), the "my videos" library (P6), and
in-UI annotation (P7?P9).
80     **Remaining = owner-gated only:** P10 (corpus promotion) + P11 (train/re-eval ?F1) +
the `L` beta follow-ups.
81     §7 tracker is the source of truth.
82     - **Single-box appliance (KheISutra Studio)** ? **? Tracked design:
`docs/LOCAL_APPLIANCE_ARCHITECTURE.md` (DRAFT, owner-gated). **
83     The `web/` SPA drives the full **ingest (Drive / bucket URI / local upload) ? index ?
pick rallies ? reel ? download**
84     loop on ONE box, behind an **edge auth gate** (engine stays `127.0.0.1`, CORS `**` is
moot). It is the single-box
85     profile of the engine's decoupled `StorageBackend` + `backend.worker` architecture.
MVP ran on the CPU droplet
86     (retired 2026-06-25) via **Gemini/stub**; **real local-GPU detection needs a GPU
host** (any VM provider / Cloud
87     Run+GPU ? owner priority ASAP, runs parallel to the UI MVP). Per-rally selection = one
screen. Owner decisions pending in §10 (edge auth,
88     retention, cost ceiling, GPU host, P5 bridge).
89     - **Alpha SPA** in `web/` (Vite + React + TS + Tailwind) ? hosts the operator console +

```

the per-rally selection UI above.

```
90      - **BYO-storage + autosync** ? connect Drive/OneDrive/GCS/S3, watch the folder, auto-prepare reels
91      (engine `PLATFORM_ARCHITECTURE.md S3`). The form's GPU/storage answers are the demand signal.
92      - **Control plane** (`control-plane/`) at beta.
93
94      ## Org & repos
95      - ? **DONE (2026-06-20): proprietary repos consolidated into the `Khelsutra` org (private). ** Transferred
96      `badminton-highlight-indexer` (engine), `sports-data-collector`, `sports-obsreport` ? `Khelsutra/` (stay private; no
97      open PRs/forks; GitHub redirects git ops so active sessions kept working; local remotes repointed + verified).
98      `rally-annotator` stays in the org **PUBLIC** (OSS MIT plugin the site links to ? do not make private). Owner is org
99      admin (full access) and can **invite a contributor** (org member / per-repo outside collaborator) anytime.
100     - ? **`khelsutra` (this site repo) ? owner step pending.** A transfer breaks the **Cloudflare Workers-Builds** git
101     connection (auto-deploy) until the Cloudflare GitHub App is installed on the `Khelsutra` org and the Worker
102     reconnected (dashboard). When ready: install the CF app on the org ? `gh api repos/avidullu/khelsutra/transfer -f
103     new_owner=Khelsutra` ? reconnect the Worker to `Khelsutra/khelsutra` ? verify a deploy.
104     - ? **Engine CI under the org** ? verify the engine's GitHub Actions run on the next PR (free orgs default Actions-on;
105     enable in org settings if a run doesn't start).
106
107     ## Convention
108     Substantial work gets a tracked doc in `docs/`; one PR per item off `origin/master`; merge to `master`
109     auto-deploys (CI must be green).
110
```

25. assistant (2026-07-03T16:41:13.446Z)

Now let me check the actual README.md to see if there's guidance on where trackers should live:

26. user (2026-07-03T16:41:13.970Z)

```
1      - [RALLY_QUALITY_RESEARCH.md](RALLY_QUALITY_RESEARCH.md) ? **? Foundation doc** for closing the local-vs-Gemini rally-detection gap: the measured gap + the code-grounded over-merge diagnosis ("shuttle moving ? rally"), an inventory of the eval/experiment harness + its gaps, the strengthened evaluation framework (segmental Edit/F1@k that make over-segmentation first-class, golden-corpus + active learning, Gemini-agreement as a *shadow* signal), an **adversarially-verified, cited external-research synthesis** (TAL/TAS boundary heads, racket-sports rally-state cues, small-data, eval rigor), a **prioritized cheap?durable roadmap with falsifiable success criteria**, and (§7) a **tracked project plan** ? work-item checklist, sequencing, the single quality **number** attributed to the effort, and a **definition of done** (project completes when all tiers land + the number is recorded). The plan that the [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) log executes against; ties together EXPERIMENT_HARNESS / MULTI_SIGNAL_FUSION_PLAN / ANALYZERS_AND_RECONCILERS / OWNED_MODEL_IMPLEMENTATION_PLAN.
2      - [RALLY_DETECTION_GAP_CLOSURE.md](RALLY_DETECTION_GAP_CLOSURE.md) ? **Tracked execution doc (2026-06-18)** for closing the remaining rally-detection accuracy gap **per video, before launch**, driven by the growing golden corpus. First ground-truth clip (GX010141) reframed it: boundaries are tight (R2/R3/R4 holds) ? the gap is **detection**: (G1) recall (local & Gemini both miss ~half on hard clips, and *different* rallies ? complementary), (G2) precision (local over-fires), (G3) **rally-END over-extension on the quick-return-after-point** (shuttle still moving ? needs a *player-state* signal). Levers (owner-gated): player-kinematics + shuttle-in-hand / sustained-invisibility end-rule with reverse-timeline backtrack, `rally_gate` Gate A wiring, complementary fusion. **Peer** of `RALLY_QUALITY_RESEARCH.md` (execution tracker, not research); §7 is the source of truth.
```

3 - PROMINENT_COURT_DETECTION.md ? **Design + experiment** (2026-06-19, default-OFF, owner-gated flip)** for the **multicourt G2 over-fire**: identify the **prominent (foreground) court** (the one covering the majority of the frame) and keep only rallies whose shuttle lives inside it ? re-defining the shuttle signal `(x,y,visible)` ? `(x,y,visible-AND-inside-prominent-court)` **before** windowing, so a background/adjacent-court shuttle never forms a rally. 2D floor polygon (MVP) ? **pseudo-3D play-volume** (refinement, so high clears aren't clipped). Born from the GX010142 rally-#1 background-court FP; reuses `fusion\_features.point\_in\_polygon` + `FusionConfig.court\_polygon`. §8 progress tracker is the source of truth.

4 - RALLY_DETECTION_QUALITY_REPORT.md ? **Technical-paper checkpoint** (2026-06-17)** of rally-detection quality: the honest measured state (over-segmentation milestone tolerance-F1 0.09170.323 at n=11, p=0.002; merge-rate 0.69470.150; first real-footage ground-truth at-par with START at Gemini parity), the task-faithful R2 metric, a detailed Gemini comparison, next steps + expected wins, areas yet to be explored, and a publishability / path-to-submission verdict. Renders to a polished PDF; the externally-shareable "toward a paper" summary of `RALLY\_QUALITY\_RESEARCH.md` + `QUALITY\_ITERATIONS.md`.

5 - EVALUATION.md ? How we measure and improve quality.

6 - EXPERIMENT_HARNESS.md ? Design (P1 SHIPPED) for A/B + shadow + SxS experimentation on rally-detection variants with **zero production regression** ? experiments are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift, rollback = flip a flag). Composes the existing `backend/eval/` machinery (`calibration`, `regression`, `metrics`); implemented in `backend/eval/experiment.py`.

7 - QUALITY_ITERATIONS.md ? **Living log** of rally-detection quality iterations (hypothesis ? what shipped ? measured/expected effect ? lesson). The reverse-chronological record; terminal-state snapshots get archived under `archives/quality-iterations/`.

8 - ANALYZERS_AND_RECONCILERS.md ? Design (owner-gated, NOT started): a **signal-fusion framework** for rally detection ? producers emit confidence-scored **signals** at three temporal **scopes** (frame detectors / window analyzers / session analyzers), a learned **scorer** weights them, and an auditable **report-first reconciliation** pass lints the decision against the evidence. The spine: **no special-casing** in the decision path (signals?scorer, never `if/else`). Worked examples: a service-fault analyzer, a warm-up/practice span detector, and consuming a per-frame shuttle-state detector. A peer to `EXPERIMENT\_HARNESS.md` and `MULTI\_SIGNAL\_FUSION\_PLAN.md` (ADR-007 modular rally layer).

9 - PERSONALIZATION_PLAN.md ? Design (owner-adopted **hybrid**; Phase-0 landing) for the **per-user personalization** feature on a generalization-first baseline: a thin per-user **tuning** layer gated by the honest small-N statistics (a **min-N promotion floor** + a **structural-guard exemption**, never a service gate) + a **contribution flywheel** (free-tier videos pool into the owner-trained baseline; "a tool that gets better the more you help it"). Records the **locked owner decisions** (identity, storage, `min\_n` semantics, Gate-A default, free-tier pooling). Built so far: per-user promotion gate (#141) + held-out-USER learning curve (#142).

10

11 - GOLDEN_REGRESSION_FIXTURES.md ? **Design** (owner-gated, NOT started): **video-free, non-attributable regression fixtures** so a clean clone runs the rally-seg suite with **no video/GPU/DB**. Two layers ? the post-detector freeze + **dual gates** (characterization + quality-floor) reusing the `backend/eval/` stack, and the **SEALED-FIXTURES** privacy filter (de-attribution-at-source, tiered public/key-gated, shuttle-only public tier). Carries a cited **IN/US/EU legal memo** (derived-data IP/privacy), a threat model, and a **§7 progress tracker + definition-of-done**. Extends [GOLDEN_DATA_SHARING.md](data_pipeline/GOLDEN_DATA_SHARING.md).

12

13 **## ? Documentation status & governance**

14 - [AUDIT_REPORT_khelsutra-guru_2026-07-02 (? COMPLETE / ARCHIVED)](archives/past_projects/AUDIT_REPORT_khelsutra-guru_2026-07-02-COMPLETED-2026-07-02.md) ? the five-repo `khelsutra-guru` workspace code audit + its folded-in **Remediation tracker** (P0?P28 all closed). **Archived 2026-07-02** once every remediation wave landed and the last residuals were closed or filed as issues (#443 CORS, #444 corpus Phase-2, #445 GCS eval; + #301/#327/#332/#384). Absorbed the prior `CODE\_CLEANUP\_AUDIT\_2026-06-24` + June-2026 hardening sweep (both archived).

15 - DOC_STATUS.md ? **living Documentation Status & Archival Register**: every foundational/tracked doc's lifecycle + archive-readiness, the **archival due-diligence checklist** (*honestly update ? then archive*), and the open doc-hygiene findings. Review here; update rows as docs complete.

16

```

17     ## Product / platform plans
18     - [ALPHA_LAUNCH_READINESS.md](ALPHA_LAUNCH_READINESS.md) ? *** Active tracked project
('DRAFT', created 2026-07-03):** cross-repo alpha readiness tracker after the golden corpus
reached 15 videos. Separates the hosted product-alpha lane (shareable upload?process?pick?reel)
from the owned-model/commercial rally-detection gate (ship target + WASB-C3 still unresolved).
$7 is the source of truth for the concrete next PR sequence: corpus portability, n=15 baseline
refresh, ablation/fusion reruns, hosted serving, alpha copy/launch review.
19     - [COMPUTE_DECOUPLED_SERVING/](COMPUTE_DECOUPLED_SERVING/README.md) ? *** Active tracked
project (2026-06-21, 'DRAFT' owner-gated):** the productionization successor to
DECOUPLED_COMPUTE (its deferred M3/M4). **One pure core (DETECT?WINDOW?STITCH) that never
branches on deployment profile**, with a 'ComputeBackend' (where it runs) + 'StorageBackend'
(what bytes) selected by **one startup config**. Ships **truly-local** (local/USB + local GPU,
in-process stitch, zero cloud) + **cloud-serving** (upload<2GB / gdrive / bucket ? a **Cloud Run
GPU job runs detect AND stitch**, metered into a per-job cost ledger) + **cpu-mock** (CI).
Includes [TESTING_STRATEGY.md](COMPUTE_DECOUPLED_SERVING/TESTING_STRATEGY.md) (profile-PARITY
proof for $0), [architecture.svg](COMPUTE_DECOUPLED_SERVING/architecture.svg), a
[runlogs/](COMPUTE_DECOUPLED_SERVING/runlogs/) posterity dir, and **ADR-014** ([lineage
ADR-009?010?011?012?013?014](archives/decisions/DECISIONS.md)). $7 M0?M10 tracker is the source
of truth.
20     - [DECOUPLED_COMPUTE/](archives/past_projects/DECOUPLED_COMPUTE/README.md) ? ?
**DELIVERED + ARCHIVED (2026-06-21).** Lifted the rally pipeline onto a **rented cloud-native
GPU** doing bucket-driven 'train' + 'infer': storage/compute seams (PRs #246?#259) + the **M3.5
native WASB runner**; **M0?M2 + M3.5 proven on a GCP L4** (M2 ? 'weights/0.1.0-l4/'). Delivered
on **GCP** ? DigitalOcean dropped for compute ([ADR-013](archives/decisions/DECISIONS.md)).
**M3/M4 (serving endpoint + per-video billing + scale-to-zero) deferred ? the "Cloud Run + GPU
serving" subproject** ([ADR-012](archives/decisions/DECISIONS.md)). Snapshot under
'docs/archives/past_projects/DECOUPLED_COMPUTE/' (incl. 'DEPLOY_REPORT_GCP_2026-06-20.md').
21     - [PACKAGING_PLAN.md](PACKAGING_PLAN.md) ? *** Active tracked project ('IN PROGRESS', 11
PRs merged 2026-06-21/22):** ship the engine as a **downloadable batteries-included Python app**
? 'pip install' ? 'indexer --app' boots a local webserver and opens the **bundled KhelSutra
SPA** single-origin, **torch-free** (LIGHT tier), cross-platform, no native component (decision
D10/O1). Built ON the decoupled-compute seam. **Supersedes the 'DESKTOP_APP_PLAN.md'
narrative.** $7 tracker = source of truth (P0a?e ?, P1a/b ?, P2-launcher/P2a/P2d/P2e ?). ***
Next:** P3 GPU validation + screencast ? P2b-install ? P2c wizard (cross-repo) ? P4
(ONNX/train/annotate = north star). Memory '[[packaging-app-plan]]' / '[[next-session-packaging-
app]]'.
22     - [DESKTOP_APP_PLAN.md](DESKTOP_APP_PLAN.md) ? *** Superseded narrative (by
'PACKAGING_PLAN.md'; D10 = self-hosted webserver, NOT a desktop app).** Original scoping doc
(2026-06-18, '[design]'): package the local pipeline as a **cross-platform desktop app** for
non-technical players ? plug in the camera's USB/SD card, click once, get back **only the
highlight reels**, fully on-device (no upload, no original-video bloat). The centerpiece is the
**de-WSL native-compute port** (run WASB natively per-OS: Linux/Windows CUDA · macOS MPS/CPU ?
drop WSL2), and its **make-or-break unknown**: is CPU acceptable on a typical *no-GPU*
enthusiast laptop (WASB is ~3.4x real-time on a laptop 2070S)? $7 tracker = **P0 compute-
feasibility spike ? P1 native detector ? P2 app shell ? P3 packaging/installers**. Realizes the
**L0 fully-local** tier of [VIDEO_LOCALITY_MODEL.md](VIDEO_LOCALITY_MODEL.md); is the
**'LocalDesktop' ComputeBackend** of [PLATFORM_ARCHITECTURE.md](PLATFORM_ARCHITECTURE.md) $3.2;
the smaller/export-clean owned model
([OWNED_MODEL_IMPLEMENTATION_PLAN.md](OWNED_MODEL_IMPLEMENTATION_PLAN.md)) is one answer to the
compute crux.
23
24     ## Practical / ongoing
25     - [CREATING_REELS.md](CREATING_REELS.md) ? *** End-user guide:** install the LIGHT-tier
app ? 'indexer --app' ? process a video ? pick rallies ? download a watermarked reel. Task-first
usage walkthrough (web UI + the API chain), with a $6 pointer to the cloud compute option. LIGHT
tier only (offline 'motion' / optional Gemini; no on-device GPU or one-click cloud). Install
mechanics ? root ['README.md'](..../README.md); packaging roadmap ?
['PACKAGING_PLAN.md'](..../PACKAGING_PLAN.md).
26     - [CONVERT_GCS_VIDEOS_TO_HIGHLIGHTS.md](CONVERT_GCS_VIDEOS_TO_HIGHLIGHTS.md) ? ***
Operator runbook:** 'gs://' bucket videos ? highlight reels on the approved **L4 Cloud Run**
(scale-to-zero, $0 idle). Copy-pasteable: build image ? create the GPU Job ? cloud-serving
control plane ? process(cloud)?query?compile?download ? prove the path
('verify_compute_path.py') ? $0 teardown. **Verified e2e 2026-06-23.** Raw deploy knobs ?
['../deploy/cloudrun/README.md'](..../deploy/cloudrun/README.md).
27     - [SETUP_NEW_MACHINE.md](SETUP_NEW_MACHINE.md) ? *** The canonical "spin up a box ?

```

ready for training & inference" runbook.** Copy-pasteable, human-followable, **verified end-to-end on DigitalOcean (2026-06-20)** ? but **compute is now GCP-only (ADR-013)**, see [`deploy/gcp/README.md`](../deploy/gcp/README.md): credentials ? provision ? bootstrap ? suite (922 passed) ? secret-free CPU inference (no GPU/keys) ? real Gemini inference ? Gen-0 training, with pointers to the GPU/native-detector + bucket paths. The `deploy/digitalocean/\*` scripts are provider-agnostic and reused on GCP.

28 - DATA_IN_GCS.md ? **? The corpus source-of-truth map: where every piece of data lives in `gs://khehsutra-rally-corpus`** (canonical layout · current contents + drift · how `backend.worker` consumes it · the still-local-only **gap** · migration plan). Read this so a session pulls data from the bucket, not from someone's `C:\`/`F:\`/Drive ? the goal is to remove the local box as a blocker. Pairs with the two runbooks below.

29 - [`deploy/gcp/nightly\_regression/README.md`](../deploy/gcp/nightly_regression/README.md) ? **? Nightly engine-regression OPERATIONS MANUAL (one-stop shop):** enable · run on demand · change any setting (email recipients/frequency/clips/config) · observe · pause/disable/remove. The full WASB-pipeline reel-count+cost+email nightly on an ephemeral GCP GPU VM (design/status: NIGHTLY_ENGINE_REGRESSION.md). The lighter CPU cached-trajectory tripwire is `scripts/nightly\_regression.py` (local Windows Scheduled Task, PR #202).

30 - TRACKNET_WSL_SETUP.md

31 - Evaluation harness and calibration drivers (in `backend/eval/`)

32
33 ## Data layer (collection · annotation · golden sets ? **no ML**)
34 - [data_pipeline/](data_pipeline/README.md) ? **the isolated data pillar**: how footage is recorded (`VIDEO\_RECORDING\_GUIDELINES`), golden sets are created + which videos (`GOLDEN\_SET\_IMPLEMENTATION` · `GOLDEN\_SET\_PHASE1\_VIDEOS` · `GOLDEN\_VIDEOS`), and data is shared/ingested (`GOLDEN\_DATA\_SHARING`). Deliberately **no segmentation/stitching/model** detail ? that's the ML pillar (top level + `COMPUTE\_DECOUPLED\_SERVING/`). The **annotation-vendor (Roora)** engagement is archived under `archives/roora-vendor/`.

35
36 ## Conventions & templates
37 - PROJECT_DOC_TEMPLATE.md ? Reusable skeleton for a **tracked project doc**: status header · decisions-locked · progress tracker (??/?) · definition-of-done · archived on completion. Standardizes the emergent pattern in `RALLY\_QUALITY\_RESEARCH.md` §7 and `GOLDEN\_REGRESSION\_FIXTURES.md`. Use for any substantial design/project work.

38
39 ## Historical documents
40 Completed plans, old research, ADRs, and checkpoints live in archives/.

41
42 - **June-2026 test-hardening sweep** ? **COMPLETED** (2026-06-14; T1?T5 + full loop + verification per #162) and **archived**: the former top-level pointer-stubs were removed 2026-07-02 (absorbed into the umbrella's *Prior audits*). Full records: archives/CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md, archives/DEFERRED_HARDENING_PLAN-COMPLETED-2026-06-14.md, archives/HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md.

43 - **CODE_CLEANUP_AUDIT (2026-06-24)** ? the prior deep engine audit (35 findings; Phase 1 complete, Phase 2 largely shipped, remaining = the `main.py` god-module split). **Absorbed into the umbrella** (see its *Prior audits (absorbed)*) and archived ? archives/past_projects/CODE_CLEANUP_AUDIT_2026-06-24-SUPERSEDED-2026-07-02.md.

44 - **KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN** ? the five-repo audit execution board, **folded into the umbrella** as its *Remediation tracker (live)* and archived ? archives/past_projects/KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN-FOLDED-2026-07-02.md.

45 - DECISION_SNAPSHOT_REGRESSION.md ? **? Tracked project (decisions D1?D11 LOCKED 2026-06-25, #383; implementation P1?P7 pending):** a deterministic **quality-neutral proof** ? freeze + compare the pipeline's **decision output** (rally segments + stitch/concat plan) off a cached trajectory, GPU-free, in CI. The behavioral proof that lets a *genuine* (non-AST-identical) refactor be confidently called quality-neutral. **Never hashes the encoded reel** (non-deterministic). Extends the existing characterization snapshot (`test\_golden\_fixtures.py`) from segmentation to the stitch plan; peer of the M4 quality-floor + GPU `nightly\_reelcount.py`. §7 tracker is the source of truth.

46

47 - ****Field-ops / customer-feedback**** (recruiting flyer · field-visit playbook · PMF) ?
[``archives/field-pmf/``](archives/field-pmf/) (archived 2026-06-21; PMF validation still an
active **business** thread ? see ``NEXT_STEPS.md``).
48 - ****Roora annotation-vendor**** engagement (labeling brief · vendor bake-off) ?
[``archives/roora-vendor/``](archives/roora-vendor/) (archived 2026-06-21).
49 - Superseded: per-video output layout ? [``archives/PER_VIDEO_OUTPUT_LAYOUT-
SUPERSEDED-2026-06-21.md``](archives/PER_VIDEO_OUTPUT_LAYOUT-SUPERSEDED-2026-06-21.md) (?
``PER_VIDEO_WORKSPACE.md``).
50
51 See ``archives/README.md`` for the umbrella policy and ``archives/past_projects/README.md``
for classification of small tech debt / starter projects.
52

27. assistant (2026-07-03T16:41:16.211Z)

Now let me analyze the key convention points. Let me check if `DOC_STATUS.md` references a
convention about where cross-repo work is tracked:

28. user (2026-07-03T16:41:17.559Z)

21:3. ****Update ALL inbound references**** ? ``docs/README.md`` index, ``CLAUDE.md``,
``docs/NEXT_STEPS.md``, and any cross-doc links ? the new path + status.
33:| ``ALPHA_LAUNCH_READINESS.md`` | ****DRAFT**** ? cross-repo alpha readiness tracker created
2026-07-03 after the golden corpus reached n=15; separates hosted product-alpha work from owned-
model/commercial rally-detection gates | No | Execute §7 in order: corpus portability + n=15
baselines, hosted alpha serving, ablation/fusion reruns, then launch GO/NO-GO |
78:``PROJECT_DOC_TEMPLATE.md`` · ``NEXT_STEPS.md`` (? refreshed 2026-06-21) · ``BACKLOG.md`` ·
``EVALUATION.md`` · ``ABLATION_LAYER.md`` · ``SETUP_NEW_MACHINE.md`` · ****`CREATING\_REELS.md`**** (? end-
user usage guide ? install ? process ? pick ? reel; LIGHT tier; added 2026-06-23) ·
****`CONVERT\_GCS\_VIDEOS\_TO\_HIGHLIGHTS.md`**** (? operator runbook ? ``gs://`` videos ? reels on L4
Cloud Run; verified e2e 2026-06-23) · ``TRACKNET_WSL_SETUP.md`` · ``README.md``. *(2026-06-21
relocation: the data-layer guides moved ? ``data_pipeline/``; the field-ops + Roora docs ?
``archives/{field-pmf,roora-vendor}/``.)
90:2. ****? RESOLVED (2026-06-21) ? `NEXT\_STEPS.md` refreshed**** with a current 2026-06-21 state +
a ****cross-track prioritized pending list****; the 2026-06-13 owned-model/quality body is kept
(clearly marked) as that track's detail.

29. assistant (2026-07-03T16:41:19.867Z)

Now let me look specifically at whether there are two competing sources of truth - `NEXT_STEPS` P0
item vs the new tracker doc:

30. user (2026-07-03T16:41:20.824Z)

0. ****? [owner-priority, 2026-06-21] MAKE IT UI-DRIVEN + ALPHA-SHAREABLE ? "a few recordings
lying around ? highlights, from a page, no SSH".**** Today the cloud flow is ****SSH/CLI-only****
(``docs/CLOUD_GPU_DRYRUN_GUIDE.md``); ****M12 does NOT fix it.**** Needs (a separate khelsutra ``web/``
SPA + deploy track, NOT an engine seam): **** (i) **** deploy the engine as a reachable ****service****
(M7 ? Cloud Run service, ? owner GCP spend for a **persistent** endpoint), **** (ii) **** the ****SPA
ingest/progress screen**** (khelsutra ``web/`` B-P2: enter/upload video ? ``/api/process`` ? poll
``/api/jobs`` ? reel; ZERO engine code), **** (iii) **** flip ****M9 edge-auth ON**** (merged #290, default-
OFF) to gate alpha tenants, **** (iv) **** a ****compute PICKER in the UI**** ("your machine vs our
cloud", D13/D15). The engine seams make it **possible**; the SPA makes it **usable**. Full record:
[[ui-driven-cloud-serving-gap]]. ****Concrete execution tracker after the n=15 corpus review:****
``docs/ALPHA_LAUNCH_READINESS.md`` §7.
0b. ****? [packaging track, IN PROGRESS ? 11 PRs merged 2026-06-21/22] Downloadable batteries-
included app ? `docs/PACKAGING\_PLAN.md`.**** The LIGHT-tier engine is bundle-ready: ``pip install``
? ``indexer --app`` serves the ****bundled SPA**** single-origin, ****torch-free****, with app-home state
/ ffmpeg resolver / DNS-rebinding guard / single-instance lock / safe DB upgrades (P0a?e, P1a/b,
P2-launcher/P2a/P2d/P2e). ****? Next (fresh focused pushes):**** **** (1) **** P3 ****GPU validation +
screencast**** (L4, ~\$0.50, DELETE-after, \$100 cap ? ``[[gcp-vm-always-delete]]``; ? native-WASB
0-rally tuned-config blocker), **** (2) **** P2b-install (uv/script + truly-local default config),
**** (3) **** P2c first-run wizard (cross-repo khelsutra SPA from ``/api/capabilities``), then P3-cloud
(Cloud Run) + P4 (ONNX/train/annotate = north star). Memory ``[[packaging-app-plan]]`` / ``[[next-`

session-packaging-app]]`.

1. ? **COMPUTE_DECOUPLED_SERVING** M0?M9 LARGELY SHIPPED** (2026-06-21, ADR-014 ratified). M1?M4 (#279/#280/#282/#283), **M5 #286, M6 #287, M8 #288, M9-auth #290**** all merged (default-OFF); the **real M7 L4 run**** validated the worker (22 rallies, VM deleted, \$0). All S8 gating Qs locked (D7?D17). **Owner residual**** (not blocking Claude's default-OFF code): counsel ToS/DPA (M9-consent live + M11 live), real GCP prices+FX (M8 calibrate), Google OAuth app (M12 real), CF team-domain/AUD + ingress-lock; real cloud verification authorized within a **\$100/?10k cap**** (D17).

2. **[owner, ongoing] Grow the golden corpus**** ? multi-court badminton + ?3 matches/sport (TT first). The single highest-leverage move; the quality track **and** full-corpus training feed on it.

31. assistant (2026-07-03T16:41:23.520Z)

Now let me check the S7 of the new ALPHA_LAUNCH_READINESS.md to understand what the source of truth claim is:

32. user (2026-07-03T16:41:24.211Z)

91 ## 7. Deliverables & progress tracker

92

93 Legend: TODO = not started, IN PROGRESS = actively being worked, DONE = shipped/verified, GATED = blocked on owner/legal/platform decision. One PR per row unless the row explicitly names a cross-repo batch.

94

95 | ID | Deliverable | Repo / owner | Depends on | Gated? | Status | PR |

96 |----|-----|-----|-----|-----|-----|----|

97 | A0 | Create this alpha readiness tracker and index it in docs | `badminton-highlight-indexer` | - | No | IN PROGRESS | #451 |

98 | A1 | Stand up alpha serving endpoint: Cloud Run/GPU or approved host, health check, edge auth ON, upload -> process -> jobs -> compile -> download verified | `khelsutra` + `badminton-highlight-indexer` | A0 | Owner spend / CF config | TODO | - |

99 | A2 | Make the n=15 golden manifest portable/reproducible; add a path/corpus smoke check that fails loudly when labels or trajectories are missing | `badminton-highlight-indexer` | A0 | No | TODO | - |

100 | A3 | Promote the 15-video source-video/MD5 records into the collector/vault path; reconcile collector's six-video source manifest with the 15-video eval corpus | `sports-data-collector` + vault | A2 | Owner upload/storage scope | TODO | - |

101 | A4 | Refresh the n=15 nightly regression baseline intentionally and record the exact command/output; keep stale-baseline warnings until reviewed | `badminton-highlight-indexer` | A2 | Reviewer sign-off | TODO | - |

102 | A5 | Recompute the heuristic n=15 floor and replace the stale `heuristic\_lovo\_n6` floor for future Gen-0 comparisons | `badminton-highlight-indexer` | A2 | No | TODO | - |

103 | A6 | Fix `backend.eval.ablation` CLI circular import, then run the R2/tolF1 ablation on the n=15 corpus | `badminton-highlight-indexer` | A2, A4 | No | TODO | - |

104 | A7 | Regenerate or upconvert full-fusion feature JSONs for the newer 9 videos so fusion/person/audio analyses run over all 15 | `badminton-highlight-indexer` | A2 | GPU/data availability | TODO | - |

105 | A8 | Re-run `fusion\_compare`, group ablation, and served contrast after A7; explicitly promote/reject person fusion, Gate A served-default, and any other R-series default flips | `badminton-highlight-indexer` | A6, A7 | Reviewer sign-off | TODO | - |

106 | A9 | Run Gen-0/TemporalMaxer shadow training against the refreshed n=15 floor; produce a non-serving artifact bundle and gate verdict | `badminton-highlight-indexer` | A5, A8 | Owner approval for M0.4 scope | TODO | - |

107 | A10 | Decide and record the alpha ship-gate target: metric, corpus slice, minimum threshold, significance rule, and what claims are allowed | Owner + `badminton-highlight-indexer` | A4, A5, A8 | Owner decision | GATED | - |

108 | A11 | Resolve or explicitly defer WASB C3 provenance for alpha messaging; if deferred, document "no commercial model sharing" copy guardrail | Owner/legal | A10 | Legal/provenance | GATED | - |

109 | A12 | Reopen or defer P10/P11 corpus flywheel: `PATCH /api/segments`, collector promotion, `train|reeval` job, and delta-F1 surface | `khelsutra` + `badminton-highlight-indexer` + collector | A1, A3 | Owner D13 reopen | GATED | - |

110 | A13 | Alpha launch review: runbook, privacy/consent copy, known limits, rollback/disable steps, and owner G0/N0-G0 | Cross-repo | A1, A4, A10, A11 | Owner launch

```

decision | TODO | - |
111
112     ## 8. Concrete next PR sequence
113
114     1. **Docs PR:** land this tracker.
115     2. **Corpus portability PR:** rebase/portable manifest tooling plus missing-file smoke
test.
116     3. **Hosted alpha PR/batch:** health/auth/upload/process/job/compile/download
verification and runbook updates.
117     4. **Baseline PR:** intentional n=15 nightly baseline refresh and heuristic floor
refresh.
118     5. **Ablation/fusion PRs:** fix ablation import, complete 15-video features, rerun
promotion/rejection decisions.
119     6. **Launch decision PR/docs update:** record ship-gate target, C3 status, alpha copy
limits, and launch checklist.
120
121     ## 9. Open questions
122
123     1. **Owner:** Is the first alpha allowed to use manual/offline corpus promotion, or
should P10/P11 be reopened before any tester wave?
124     2. **Owner/platform:** Which hosted compute target is the next stand-up: Cloud Run GPU,
GCE L4 VM, or another approved GPU host?
125     3. **Owner/legal:** What is the acceptable alpha wording while WASB C3 remains
unresolved?
126     4. **Engineering:** Should the portable corpus manifest live as a generated artifact, a
checked-in redacted manifest, or both?
127     5. **Engineering:** Should n=15 feature regeneration target full fusion schema for every
clip, or should the eval loaders support mixed schemas with explicit capability tags?
128
129     ## 10. Definition of done
130
131     - [ ] Invited tester can complete the hosted record/upload -> reel download CUJ without
SSH.
132     - [ ] Alpha service is gated behind the chosen auth posture and has a documented
disable/rollback path.
133     - [ ] n=15 corpus is portable enough that manifest-driven quality commands do not rely
on stale local absolute paths.
134     - [ ] n=15 nightly baseline and heuristic floor are refreshed with recorded commands.
135     - [ ] R2/tolF1 gate is promoted or rejected with a recorded ablation result.
136     - [ ] Fusion/person/Gate-A served-default decisions are explicitly promoted or rejected
on the 15-video corpus.
137     - [ ] Ship-gate target and allowed alpha claims are recorded.
138     - [ ] WASB C3 status is recorded and reflected in alpha messaging.
139     - [ ] P10/P11 are either reopened with a committed sequence or explicitly deferred for
alpha-plus.
140     - [ ] Owner records GO/NO-GO for the first alpha wave.

```

33. assistant (2026-07-03T16:41:51.200Z)

Perfect. Now I have all the information I need to do the altitude analysis. Let me compile my findings based on the conventions in DOC_STATUS.md.

Based on my review of the repo's documentation conventions, here are the altitude-related findings for PR #451 (ALPHA_LAUNCH_READINESS.md):

```

**Finding 1: Two competing sources of truth for the same P0 work (NEXT_STEPS.md vs new tracker
doc)**
- **File:** docs/NEXT_STEPS.md, line 11
- **Summary:** NEXT_STEPS.md P0 item (line 11: "MAKE IT UI-DRIVEN + ALPHA-SHAREABLE") claims to
be the "Full record" for UI-driven cloud serving work, while ALPHA_LAUNCH_READINESS.md §7 now
claims to be "the source of truth for the concrete next PR sequence." DOC_STATUS.md §2.1
requires updating "ALL inbound references" when a doc's lifecycle changes, but NEXT_STEPS.md P0
remains as a live, parallel reference without demotion or subordination language.
- **Failure scenario:** A contributor updating the P0 item in NEXT_STEPS.md will not auto-sync
with the tracker in ALPHA_LAUNCH_READINESS.md §7, causing split information and stale decision

```

records.

****Finding 2: ADR/decision records locked in a DRAFT tracker doc (not in formal ADR log)****

- ****File:**** docs/ALPHA_LAUNCH_READINESS.md, lines 27-35 (§2 "Decisions locked")
- ****Summary:**** The new doc locks five decisions (D1?D5) inside a DRAFT tracker document, bypassing the ADR mechanism documented in DOC_STATUS.md §3f and exemplified in DECISION_SNAPSHOT_REGRESSION.md §2. Per DOC_STATUS.md §2.1 archival checklist, "decisions locked" should flow through the ADR log (`docs/archives/decisions/DECISIONS.md`). The existing convention shows decisions like "D7?D17" formalized as ADR-014 ratified (2026-06-21, per DOC_STATUS.md §6); D1?D5 here are tactical alpha decisions but are embedded in the DRAFT doc itself.
- ****Failure scenario:**** If the alpha tracker is deferred or rejected, these decisions (especially D1 "Product alpha and owned-model launch are separate decisions") will not appear in the decision lineage, making future model-launch decisions unable to reference or audit the prior reasoning.

****Finding 3: Cross-repo tracker altitude unclear (sibling repos have separate NEXT_STEPS)****

- ****File:**** docs/ALPHA_LAUNCH_READINESS.md, lines 5-6 and §7 cross-repo items (A1, A3, A12)
- ****Summary:**** DOC_STATUS.md does not specify whether cross-repo trackers should live in one repo's docs/ or in a shared location. The sibling `khelsutra` repo has its own NEXT_STEPS.md and tracked projects (e.g., `STUDIO\_MEDIA\_LIFECYCLE.md`). ALPHA_LAUNCH_READINESS.md in `badminton-highlight-indexer` now tracks work in `khelsutra`, `sports-data-collector`, and cross-repo deliverables (A1, A3, A12). The khelsutra NEXT_STEPS.md (read above) has no back-reference to this tracker.
- ****Failure scenario:**** The khelsutra team updating `STUDIO\_MEDIA\_LIFECYCLE.md` P10/P11 (corpus flywheel; mentioned in ALPHA_LAUNCH_READINESS.md A12) will not see the alpha readiness dependency, and the two repos will track the same work asynchronously.

****Finding 4: "Source of truth" claim conflicts with NEXT_STEPS.md language (not demoted)****

- ****File:**** docs/ALPHA_LAUNCH_READINESS.md, line 5 and §7 header (line 91-93)
- ****Summary:**** The doc's header (line 5) states "Section 7 is the progress tracker and source of truth," but NEXT_STEPS.md P0 (line 11, updated in this PR) still says "Full record: [[ui-driven-cloud-serving-gap]]. ****Concrete execution tracker after the n=15 corpus review:**** `docs/ALPHA\_LAUNCH\_READINESS.md` §7" ? it acknowledges the new tracker exists but doesn't demote its own P0 claim. Per DOC_STATUS.md §2, inbound references should be "repointed" and hierarchy clarified; instead, both docs now claim "source of truth" for overlapping work.
- ****Failure scenario:**** A PM updating NEXT_STEPS.md P0 (the historically-established entry point) will not realize it is subordinate to the tracker, causing updates to miss the detailed A1?A13 rows.

****Finding 5: Doc-status register entry lacks clarity on source-of-truth subordination****

- ****File:**** docs/DOC_STATUS.md, line 33 (the new row for ALPHA_LAUNCH_READINESS.md)
- ****Summary:**** The added row says this is a "cross-repo alpha readiness tracker created 2026-07-03" and "separates hosted product-alpha work from owned-model/commercial rally-detection gates," but the "Next action" column says "Execute §7 in order" without stating that NEXT_STEPS.md P0 should now be demoted to a summary/reference, not a live tracker. DOC_STATUS.md itself (§2.1) states "flip its row here" when something moves lifecycle, and §2.3 requires "Update ALL inbound references" ? but this PR updates the DOC_STATUS row and adds the README.md entry without demoting NEXT_STEPS.md P0.
- ****Failure scenario:**** A reviewer reading DOC_STATUS.md will not understand which doc owns the P0 UI-driven work going forward, and the register becomes inconsistent with its own §2 archival checklist.

****Findings 1, 4, 5 address the core issue from (3) in the review brief:**** the doc is "the source of truth" while NEXT_STEPS.md P0 remains live, creating two competing sources of truth for the same P0 work ? and DOC_STATUS.md conventions are not followed to resolve this (demotion, cross-link clarity, inbound-reference repointing).

****Finding 3 is the cross-repo altitude issue:**** an explicit repo-coordination convention is missing from DOC_STATUS.md for where multi-repo work is tracked.

****Finding 2 is the decision-recording issue:**** major decisions (D1?D5) bypass the ADR log, which per the repo's own convention should be the single lineage for locked decisions.

```json

```

[
 {
 "file": "docs/NEXT_STEPS.md",
 "line": 11,
 "summary": "NEXT_STEPS.md P0 and ALPHA_LAUNCH_READINESS.md §7 both claim source-of-truth status for the same UI-driven alpha work; DOC_STATUS §2.1 requires demoting inbound references when a doc's lifecycle changes.",
 "failure_scenario": "A contributor updates NEXT_STEPS P0 without seeing ALPHA_LAUNCH_READINESS §7 is now the tracker, causing split information and stale PRs."
 },
 {
 "file": "docs/ALPHA_LAUNCH_READINESS.md",
 "line": 27,
 "summary": "Five decisions (D1?D5) are locked inside this DRAFT tracker doc, bypassing the ADR log (docs/archives/decisions/DECISIONS.md). Per DOC_STATUS §2.1 and DECISION_SNAPSHOT_REGRESSION.md §2 precedent, major decisions should flow through the formal ADR mechanism.",
 "failure_scenario": "If the alpha tracker is deferred/rejected, these decisions (esp. D1 'Product alpha and owned-model launch separate') will not appear in the decision lineage, breaking future model-launch audits."
 },
 {
 "file": "docs/ALPHA_LAUNCH_READINESS.md",
 "line": 5,
 "summary": "Doc claims §7 is 'source of truth' while sibling repos (khelsutra, sports-data-collector) have separate NEXT_STEPS.md and tracked docs. DOC_STATUS.md provides no convention for where cross-repo trackers should live or how to sync sibling NEXT_STEPS files.",
 "failure_scenario": "The khelsutra team updating STUDIO_MEDIA_LIFECYCLE.md P10/P11 (corpus flywheel, dep. A12) won't see the alpha readiness dependency, causing asynchronous tracking across repos."
 },
 {
 "file": "docs/DOC_STATUS.md",
 "line": 33,
 "summary": "The new ALPHA_LAUNCH_READINESS row does not state that NEXT_STEPS.md P0 should now be demoted to a reference. Per §2.1, inbound references must be 'updated' and lifecycle clarified, but the row lacks this subordination language.",
 "failure_scenario": "A reviewer reading DOC_STATUS.md cannot determine which doc owns the P0 work going forward; the register violates its own §2.1 'update ALL inbound references' rule."
 },
 {
 "file": "docs/ALPHA_LAUNCH_READINESS.md",
 "line": 5,
 "summary": "The doc is indexed in README.md and DOC_STATUS.md as a 'cross-repo' tracker, but the DOC_STATUS §3a register entry (line 33) treats it as a single-repo project. The 'Repo / owner' column in §7 lists multiple repos (A1, A3, A12), conflicting with the register classification.",
 "failure_scenario": "A contributor looking at DOC_STATUS §3a (which classifies by single-repo project ownership) will misunderstand ALPHA_LAUNCH_READINESS as owned by badminton-highlight-indexer alone, missing khelsutra and collector dependencies."
 }
]
...

```