

# claude-windows / 6790052a-f068-488a-9cc6-95b0c547eeb1

## Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\6790052a-f068-488a-9cc6-95b0c547eeb1.jsonl`
- SHA-256: `139302682402540d82026fbf49be5832e0971ef10dff5016ea9c6ac525708cba`
- Source modified: `2026-07-03T17:02:47+00:00`
- Imported at: `2026-07-05T16:48:17+00:00`
- project: `C--Users-avidu-Projects-badminton-highlight-indexer`
- session_id: `6790052a-f068-488a-9cc6-95b0c547eeb1`
```

## Transcript

### 1. user (2026-06-25T16:41:52.300Z)

Hello. Can you pls sync this repo to remote head and help in reviewing <https://github.com/avidullu/muneem/pull/80>

### 2. assistant (2026-06-25T16:42:03.133Z)

I'll start by syncing the primary repo, then look at the muneem repo and PR #80. Let me gather state in parallel.

### 3. user (2026-06-25T16:42:10.555Z)

```
=== badminton remote ===
origin https://github.com/Khelsutra/badminton-highlight-indexer.git (fetch)
origin https://github.com/Khelsutra/badminton-highlight-indexer.git (push)
=== status ===
```

#### master...origin/master

```
=== fetch+ff ===
From https://github.com/Khelsutra/badminton-highlight-indexer
* [new branch]      o1-p11/gemini-remote-file-gc -> origin/o1-p11/gemini-remote-file-gc
Already up to date.
```

### 4. user (2026-06-25T16:42:10.965Z)

```
=== muneem remote ===
origin https://github.com/avidullu/muneem.git (fetch)
origin https://github.com/avidullu/muneem.git (push)
=== status ===
```

#### refactor/split-cli...origin/refactor/split-cli

```
=== recent ===
ec091d5 refactor: split cli.py (~613 lines) into cli/ subpackage
71782fd Merge branch 'refactor/move-experimental-ingestion'
4bfbb68 Merge branch 'refactor/safe-amount-rename'
0948f14 Merge branch 'docs/code-health-report'
88c3c39 Merge branch 'docs/refresh-current-state'
```

### 5. user (2026-06-25T16:42:12.473Z)

```
{"additions":656,"author":{"id":"MDQ6VXNlcjE2NzQ4Nzc=","is_bot":false,"login":"avidullu","name":"Avi"},"baseRefName":"main","body":"Splits the monolithic `cli.py` (~613 lines) into a `cli/` subpackage with focused modules.\n\n### Before\n```\nsrc/muneem/cli.py  (613 lines ? all commands in one file)\n```\n\n### After\n```\nsrc/muneem/cli.py      (9 lines ? thin re-export)\nsrc/muneem/cli/\n  __init__.py         (70 lines ? app assembly, info command, registration)\n  _shared.py          (40 lines ? _ledger, BANK_DISPATCH, _version_callback)\n  ingest.py           (130 lines ? ingest +
```

```
_ingest_bank_statement)\n??? offers.py          (35 lines ? offers_app + list command)\n???
txns.py          (75 lines ? txns_app + list command)\n??? db.py          (200
lines ? db_app + all encryption lifecycle commands)\n??? eval.py          (75 lines ?
eval_statements + _eval_statement helper)\n```\n\n### Key design decisions\n- **Explicit
registration**: Commands registered via `app.command()(fn)` instead of decorators, avoiding
circular import issues\n- **`_shared.py`**: Isolates helpers (`_ledger`, `BANK_DISPATCH`) needed
by both `__init__.py` and sub-modules\n- **Lazy imports**: `extract()` imported inside functions
to preserve test monkeypatching\n- **Entry point unchanged**: `pyproject.toml` still points to
`muneem.cli:app` ? works through the thin re-export\n\n### Verification\n- ? `ruff check .` ?
clean\n- ? `ruff format --check .` ? clean\n- ? `mypy` ? 41 source files, no issues\n- ?
`pytest` ? 478 passed\n- ? `muneem --version` ? works\n- ? `muneem info` ? works\n\nFixes
recommendation #5 from the code-health report (PR #74). Also updates the report tracker to mark
all ? high-priority items as
done.", "changedFiles": 9, "deletions": 615, "headRefName": "refactor/split-
cli", "mergeable": "MERGEABLE", "number": 80, "reviewDecision": "", "state": "OPEN", "title": "refactor:
split cli.py (~613 lines) into cli/
subpackage", "url": "https://github.com/avidullu/muneem/pull/80"}
```

## 6. assistant (2026-06-25T16:42:21.927Z)

Primary repo is synced (already up to date; one new remote branch fetched). Now let me sync the muneem PR branch and check CI + the diff.

## 7. user (2026-06-25T16:42:26.499Z)

Already up to date.

=== HEAD ===

ec091d5 refactor: split cli.py (~613 lines) into cli/ subpackage

=== CI checks ===

```
test (ubuntu-latest, 3.11)      pass    30s
https://github.com/avidullu/muneem/actions/runs/28185700568/job/83486855874
test (ubuntu-latest, 3.12)      pass    53s
https://github.com/avidullu/muneem/actions/runs/28185700568/job/83486855859
test (ubuntu-latest, 3.13)      pass    38s
https://github.com/avidullu/muneem/actions/runs/28185700568/job/83486855907
test (windows-latest, 3.13)     pass    1m22s
https://github.com/avidullu/muneem/actions/runs/28185700568/job/83486855881
ocr      skipping              0
https://github.com/avidullu/muneem/actions/runs/28185700568/job/83486856842
test (ubuntu-latest, 3.11)      pass    36s
https://github.com/avidullu/muneem/actions/runs/28185690364/job/83486823098
test (ubuntu-latest, 3.12)      pass    50s
https://github.com/avidullu/muneem/actions/runs/28185690364/job/83486823197
test (ubuntu-latest, 3.13)      pass    36s
https://github.com/avidullu/muneem/actions/runs/28185690364/job/83486823145
test (windows-latest, 3.13)     pass    2m8s
https://github.com/avidullu/muneem/actions/runs/28185690364/job/83486823122
ocr      skipping              0
https://github.com/avidullu/muneem/actions/runs/28185690364/job/83486823851
```

## 8. user (2026-06-25T16:42:28.367Z)

=== diff stat vs main ===

```
docs/code-health-2026-06-25.md | 15 +-
src/muneem/cli.py               | 612 +-----
src/muneem/cli/__init__.py      | 69 +++++
src/muneem/cli/_shared.py       | 37 +++
src/muneem/cli/db.py            | 206 ++++++
src/muneem/cli/eval.py          | 83 +++++
src/muneem/cli/ingest.py        | 143 ++++++
src/muneem/cli/offers.py        | 37 +++
src/muneem/cli/txns.py          | 69 +++++
9 files changed, 656 insertions(+), 615 deletions(-)
```

```

=== project files ===
total 149
drwxr-xr-x 1 avidu 197609    0 Jun 25 22:03 .
drwxr-xr-x 1 avidu 197609    0 Jun 25 17:33 ..
drwxr-xr-x 1 avidu 197609    0 Jun 11 00:11 .claude
-rw-r--r-- 1 avidu 197609 53248 Jun 10 22:44 .coverage
drwxr-xr-x 1 avidu 197609    0 Jun 25 22:12 .git
-rw-r--r-- 1 avidu 197609   286 May 30 23:27 .gitattributes
drwxr-xr-x 1 avidu 197609    0 Jun 10 23:14 .github
drwxr-xr-x 1 avidu 197609    0 May 30 23:27 .github
-rw-r--r-- 1 avidu 197609  1579 Jun 10 23:34 .gitignore
drwxr-xr-x 1 avidu 197609    0 Jun  2 23:10 .mypy_cache
drwxr-xr-x 1 avidu 197609    0 May 30 22:57 .pytest_cache
drwxr-xr-x 1 avidu 197609    0 May 30 22:57 .ruff_cache
drwxr-xr-x 1 avidu 197609    0 May 30 22:57 .venv
-rw-r--r-- 1 avidu 197609  4273 May 30 21:54 CLAUDE.md
-rw-r--r-- 1 avidu 197609 22509 Jun 25 22:03 DESIGN.md
-rw-r--r-- 1 avidu 197609  1184 May 31 20:41 Makefile
-rw-r--r-- 1 avidu 197609  4040 Jun 10 23:34 README.md
-rw-r--r-- 1 avidu 197609  1929 Jun 25 21:55 constraints.txt
drwxr-xr-x 1 avidu 197609    0 Jun 25 22:09 docs
-rw-r--r-- 1 avidu 197609  2578 Jun 25 22:03 pyproject.toml
drwxr-xr-x 1 avidu 197609    0 May 30 22:57 src
drwxr-xr-x 1 avidu 197609    0 Jun 10 23:14 testdata
drwxr-xr-x 1 avidu 197609    0 Jun 10 23:46 tests
drwxr-xr-x 1 avidu 197609    0 Jun 25 22:03 tools

```

```

=== CLAUDE.md? ===
CURRENT_STATE.md
DEPENDENCIES.md
ROADMAP.md
TESTING.md
code-health-2026-06-25.md
encryption-at-rest.md
feedback-antigravity-response.md
feedbackJune2026.md
ingesting-statements.md
onboarding-a-bank.md
parser-architecture.md
prior-art-email-mining.md
proposals
recovery.md
research-open-banking-india.md
schema.md

```

## 9. assistant (2026-06-25T16:42:37.583Z)

CI is green across all platforms (Ubuntu 3.11/3.12/3.13 + Windows 3.13), OCR skipped. This is a pure file-split refactor. Let me read the project conventions and do a careful semantic-preservation review of the diff.

## 10. user (2026-06-25T16:42:38.069Z)

```

1      # CLAUDE.md ? muneem
2
3      You are working on muneem, a personal-finance bookkeeper that ingests financial
4      documents from the user's email and builds a local consolidated ledger.
5
6      ## Read these before doing anything substantive
7
8      1. [ `DESIGN.md` ](./DESIGN.md) ? the full design: vision, five-stage pipeline, security
9      posture, open decisions, recommended first slice.
10     2. [ `docs/prior-art-email-mining.md` ](./docs/prior-art-email-mining.md) ? what the
11     predecessor project ( `~/Projects/Email-Mining` ) tried, what worked, what to inherit, what to
12     drop.
13     3. [ `testdata/promo-emails/README.md` ](./testdata/promo-emails/README.md) ? what's in

```

the regression corpus and where it came from.

```
10
11     ## Project context
12
13     - **Single-user, personal.** This is the user's own tool for the user's own financial
data. Not a multi-tenant product. Design decisions optimize for the user's needs, not for
hypothetical others.
14     - **User location: Bangalore, India.** Expect Indian banks (HDFC, ICICI, SBI, Axis,
Kotak, etc.), Indian credit cards, Indian brokers (Zerodha, Groww, Kuvera, etc.), Indian
merchants (Swiggy, Zomato, Flipkart, Amazon.in, BigBasket, Blinkit, etc.). Password conventions
on Indian bank/CC PDFs typically derive from name + DOB or PAN + DDMM ? design accordingly.
15     - **Greenfield.** No application code exists yet. Don't assume a stack ? it hasn't been
chosen. See `DESIGN.md` § Open Decisions.
16
17     ## Standing rules (highest priority ? these override default behavior)
18
19     ### Security
20
21     This project aggregates the user's **entire financial life** into one local store. That
is an extremely high-value target.
22
23     - **Treat all statement contents as sensitive by default.** Bank balances, card numbers,
transaction lines, holdings, PII.
24     - **Never send sensitive document contents to a cloud LLM API** (OpenAI, Anthropic,
Gemini, etc.) without explicit per-document-category approval discussed with the user in the
current session. Default to local extraction (Tesseract OCR, local LLMs via Ollama,
deterministic parsers).
25     - The corpus in `testdata/promo-emails/` is **non-sensitive** (promo emails only, no
account numbers) ? safe to send to cloud LLMs if needed for development.
26     - **Never commit:** real bank/CC/investment PDFs; the parsed DB file; OAuth tokens;
`.env` files; any password rules, password lists, or example passwords; logs containing
extracted statement text.
27     - The `.gitignore` already blocks the obvious paths; do not weaken it.
28
29     ### Decisions
30
31     There are **open architectural decisions** parked for explicit human discussion (see
`DESIGN.md` § Open Decisions). **Do not silently pick** any of them by writing code that locks
in an answer. The short list:
32
33     1. Language: Go-only, Python-only, or Go-ingest + Python-parse
34     2. Password sourcing: per-sender rule template, candidate dictionary, or interactive
prompt
35     3. LLM locality per doc category (when local OCR/parsers aren't enough)
36     4. First vertical slice scope
37
38     If the user asks you to start coding without resolving these, surface the relevant open
decisions first and ask for direction.
39
40     ### Working conventions
41
42     - Ask before introducing new external services (paid APIs, cloud vendors, hosted
services).
43     - Prefer batteries-included stdlib + well-known OSS over novel dependencies.
44     - Run all OCR / LLM work **locally by default**.
45     - Statement-handling code paths should be testable on the non-sensitive `testdata/promo-
emails/` corpus where possible.
46     - Follow attribution discipline: never invent a source, a discussion that didn't happen,
or a fact you didn't verify. If you're inferring, say so.
47
48     ## Vision continuity
49
50     muneem is the successor to `~/Projects/Email-Mining`, a coupon-finder prototype the user
built earlier. Same core insight ("there's structured value in my inbox I'm forgetting ? extract
it"); broader scope (full financial picture, not just deals). The deals/coupons feature from the
```

predecessor survives as **one of several** extraction targets in muneem. See `docs/prior-art-email-mining.md` for details and what to carry forward.

```
51
52     ## When in doubt
53
54     Ask. The user (Avi) is actively driving the design. This project is at the stage where
    shape and scope decisions matter more than code velocity. A clarifying question is almost always
    better than a guessed implementation.
55
```

## 11. user (2026-06-25T16:42:39.359Z)

```
=== ORIGINAL cli.py (from origin/main) ===
"""The ``muneem`` command-line interface.

``ingest`` parses a PDF into the local store (promo, or ``--doc-type bank_statement`` with
automatic password unlock); ``offers list`` / ``txns list <bank>|all`` read back what was
found (listings tag anything not cleanly reconciled); ``eval <folder>`` reports reconcile
rates; ``db ?`` covers the encryption lifecycle (status/encrypt/rotate-key/backup/
setup-recovery/recover). ``--ocr`` adds local image OCR for codes rasterized into banners
(approximate; flagged for verification). ``fetch`` (Gmail pull) lands with Phase F.
"""

from __future__ import annotations

from contextlib import contextmanager
from pathlib import Path

import typer

from muneem import __version__
from muneem import db as dbmod
from muneem.config import load_config
from muneem.log import get_logger
from muneem.redaction import redact_filename

_log = get_logger(__name__)

app = typer.Typer(
    add_completion=False,
    help="muneem ? a local-first personal-finance bookkeeper.",
    no_args_is_help=True,
)
offers_app = typer.Typer(no_args_is_help=True, help="Work with extracted offers.")
app.add_typer(offers_app, name="offers")

txns_app = typer.Typer(no_args_is_help=True, help="Work with parsed transactions.")
app.add_typer(txns_app, name="txns")

db_app = typer.Typer(no_args_is_help=True, help="Manage the local encrypted ledger.")
app.add_typer(db_app, name="db")

@contextmanager
def _ledger(db_path: Path):
    """Open the encrypted ledger, turning an un-migrated *plaintext* file into a clean CLI
    error (pointing at ``muneem db encrypt``) rather than a raw traceback."""
    try:
        with dbmod.connect_encrypted(db_path) as conn:
            yield conn
    except dbmod.EncryptedDatabaseError as exc:
        typer.echo(str(exc), err=True)
        raise typer.Exit(1) from exc
```

## Dispatch table for supported banks (extensible as new parsers land).

```
BANK_DISPATCH = {
    "axis": (dbmod.list_axis_transactions, "debit", "credit"),
    "axis_cc": (dbmod.list_axis_cc_transactions, "amount", "amount"), # direction from "type"
    "au": (dbmod.list_au_transactions, "debit", "credit"),
    "hdfc": (dbmod.list_hdfc_transactions, "withdrawal", "deposit"),
}

def _version_callback(value: bool) -> None:
    if value:
        typer.echo(f"muneem {_version__}")
        raise typer.Exit()

@app.callback()
def main(
    version: bool = typer.Option(
        False,
        "--version",
        callback=_version_callback,
        is_eager=True,
        help="Show the muneem version and exit.",
    ),
) -> None:
    """muneem command-line interface."""

@app.command()
def info() -> None:
    """Show resolved local paths (no secrets, no sensitive data)."""
    cfg = load_config()
    present = "present" if cfg.config_file.is_file() else "absent"
    typer.echo(f"muneem {_version__}")
    typer.echo(f"config dir : {cfg.config_dir}")
    typer.echo(f"data dir   : {cfg.data_dir}")
    typer.echo(f"db path    : {cfg.db_path}")
    typer.echo(f"staging    : {cfg.staging_dir}")
    typer.echo(f"config file: {cfg.config_file} ({present})")

def _ingest_bank_statement(pdf: Path, db_path: Path, doc_type: str, *, interactive: bool) -> None:
    """Unlock (if encrypted) ? extract ? select parser ? store reconciled transactions.

    The hybrid password unlock (cached/configured candidates ? prompt) runs *before* extraction,
    since the text needed to pick a parser is itself encrypted. The resolved password is
    forwarded
    to extraction and the parser, used in-memory only ? never logged.
    """
    from muneem import db as dbmod
    from muneem.extract import extract
    from muneem.parsers import get_parser
    from muneem.statement_unlock import resolve_password
    from muneem.unlock import is_encrypted

    password: str | None = ""
    if is_encrypted(pdf):
        password = resolve_password(pdf, interactive=interactive)
        if password is None:
            hint = " (run without --no-prompt to enter it)" if not interactive else ""
            typer.echo(f"Could not unlock {pdf.name}: no working password{hint}.", err=True)
            raise typer.Exit(1)

    ex = extract(pdf, password=password or "") # resolved above; `or ""` keeps the type str
```

```

parser = get_parser(ex.text)
if parser is None:
    typer.echo(f"No parser recognizes {pdf.name} as a supported bank statement.", err=True)
    raise typer.Exit(1)

with _ledger(db_path) as conn:
    dbmod.init_db(conn)
    existing = dbmod.get_document_by_sha(conn, ex.sha256)
    if existing is not None and existing["status"] == "reconciled":
        typer.echo(f"already ingested (doc {existing['id']}): {pdf.name}")
        return
    if existing is not None:
        # sha256 dedup must not entomb a failure: a document the oracle did NOT endorse
        # gets a fresh parse (a parser fix deserves a second chance). Same document id,
        # clean row slate; provenance upserts on the re-parse.
        doc_id = int(existing["id"])
        dbmod.clear_document_rows(conn, doc_id)
        typer.echo(
            f"re-parsing previously {existing['status']} document (doc {doc_id}):
{pdf.name}"
        )
    else:
        doc_id = dbmod.insert_document(
            conn,
            source="file",
            external_id=pdf.name,
            doc_type=doc_type,
            sha256=ex.sha256,
            path=str(pdf.resolve()),
            status="parsed",
        )
        count = parser.parse_and_insert(pdf, conn, doc_id, password=password)
        doc = dbmod.get_document_by_sha(conn, ex.sha256)
        status = doc["status"] if doc else "parsed"

typer.echo(
    f"ingested {pdf.name} (doc {doc_id}); parsed {count} transaction(s) "
    f"via {parser.__class__.__name__} [{status}]"
)

def _eval_statement(pdf: Path, *, interactive: bool) -> tuple[str, int, str | None]:
    """Run one statement through the full parse pipeline into an in-memory DB.

    Returns ``(outcome, row_count, parser_name)`` where ``outcome`` is the document status
    (``reconciled`` / ``unreconciled``) or a failure tag (``locked`` / ``no-parser`` /
    ``error``).
    Nothing is persisted and no statement value is surfaced ? only the verdict and counts.
    """
    from muneem import db as dbmod
    from muneem.extract import extract
    from muneem.parsers import get_parser
    from muneem.statement_unlock import resolve_password
    from muneem.unlock import is_encrypted

    password = ""
    if is_encrypted(pdf):
        password = resolve_password(pdf, interactive=interactive) or ""
        if not password:
            return ("locked", 0, None)
    try:
        ex = extract(pdf, password=password)
        parser = get_parser(ex.text)
    except Exception as exc:
        # eval must not crash on one bad file; record the redacted reason
        _log.warning(

```

```

        "eval: extract/route failed for %s (%s)", redact_filename(pdf), type(exc).__name__
    )
    return ("error", 0, None)
if parser is None:
    return ("no-parser", 0, None)
name = parser.__class__.__name__
try:
    with dbmod.connect(":memory:") as conn:
        dbmod.init_db(conn)
        doc_id = dbmod.insert_document(
            conn, source="file", external_id=pdf.name, sha256=ex.sha256, status="parsed"
        )
        count = parser.parse_and_insert(pdf, conn, doc_id, password=password)
        doc = dbmod.get_document_by_sha(conn, ex.sha256)
        status = doc["status"] if doc else "parsed"
except Exception as exc:
    _log.warning(
        "eval: parse/store failed for %s (%s)", redact_filename(pdf), type(exc).__name__
    )
    return ("error", 0, name)
return (status, count, name)

@app.command("eval")
def eval_statements(
    folder: Path = typer.Argument(
        ...,
        exists=True,
        file_okay=False,
        readable=True,
        help="A LOCAL folder of statement PDFs to evaluate (e.g. your seed or golden set).",
    ),
    prompt: bool = typer.Option(
        False,
        "--prompt/--no-prompt",
        help="Prompt for a password when cached/configured candidates don't unlock a file.",
    ),
) -> None:
    """Evaluate the parsers against a folder of statements and report the reconcile rate.

    The seed?golden onboarding loop (see docs/onboarding-a-bank.md): point it at a few *seed*
    statements while tuning a profile, then at a *golden* set to confirm. Everything runs
    locally in
    an in-memory DB; only counts / statuses / parser names print ? never statement values.
    """
    pdfs = sorted(folder.glob("*.pdf"))
    if not pdfs:
        typer.echo(f"no .pdf files in {folder}")
        return
    reconciled = 0
    for pdf in pdfs:
        outcome, count, name = _eval_statement(pdf, interactive=prompt)
        reconciled += outcome == "reconciled"
        typer.echo(f" {outcome:12} {count:4d} row(s) {name or '?':20} {pdf.name}")
    typer.echo(f"\n{reconciled}/{len(pdfs)} reconciled")

@app.command()
def ingest(
    pdf: Path = typer.Argument(
        ..., exists=True, dir_okay=False, readable=True, help="Path to a PDF to ingest."
    ),
    doc_type: str = typer.Option("promo", help="Document category."),
    db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data dir)."),
    ocr: bool = typer.Option(

```



```

        False, "--ocr", help="Also OCR images locally (Tesseract) for codes in banners."
    ),
    no_prompt: bool = typer.Option(
        False,
        "--no-prompt",
        help="Never prompt for a statement password; use only cached/configured candidates.",
    ),
) -> None:
    """Store a PDF's offers (promo) or transactions (bank_statement) in the local DB."""
    from muneem.extract import extract
    from muneem.offers import extract_offer

    db_path = db or load_config().db_path

    if doc_type == "bank_statement":
        _ingest_bank_statement(pdf, db_path, doc_type, interactive=not no_prompt)
        return

    from muneem import db as dbmod

    ex = extract(pdf)
    text_fields = extract_offer(ex.text)
    # (code, source) pairs; text-layer codes are exact, OCR codes are approximate.
    found: list[tuple[str, str]] = [(c, "text") for c in text_fields.codes]
    expiry_raw, expiry_type = text_fields.expiry_raw, text_fields.expiry_type

    if ocr:
        from muneem.ocr import ocr_pdf, tesseract_available

        if not tesseract_available():
            typer.echo("warning: --ocr requested but Tesseract not found; skipping OCR.",
err=True)
        else:
            ocr_fields = extract_offer(ocr_pdf(pdf))
            known = {c for c, _ in found}
            for code in ocr_fields.codes:
                if code not in known:
                    found.append((code, "ocr"))
                    known.add(code)
            if expiry_raw is None:
                expiry_raw, expiry_type = ocr_fields.expiry_raw, ocr_fields.expiry_type

    with _ledger(db_path) as conn:
        dbmod.init_db(conn)
        existing = dbmod.get_document_by_sha(conn, ex.sha256)
        if existing is not None:
            typer.echo(f"already ingested (doc {existing['id']}): {pdf.name}")
            return
        doc_id = dbmod.insert_document(
            conn,
            source="file",
            external_id=pdf.name,
            doc_type=doc_type,
            sha256=ex.sha256,
            path=str(pdf.resolve()),
            status="parsed",
        )
    for code, src in found:
        dbmod.insert_offer(
            conn,
            document_id=doc_id,
            promo_code=code,
            expiry_raw=expiry_raw,
            expiry_type=expiry_type,
            source=src,

```

```

        source_text=ex.text[:2000],
    )

    summary = ", ".join(f"{c} [{s}]" for c, s in found) if found else "none"
    typer.echo(f"ingested {pdf.name} (doc {doc_id}); {len(found)} offer(s): {summary}")

@offers_app.command("list")
def offers_list(
    db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data dir)."),
) -> None:
    """List extracted offers."""
    from muneem import db as dbmod

    db_path = db or load_config().db_path
    if not Path(db_path).exists():
        typer.echo("no offers yet (database not created).")
        return

    with _ledger(db_path) as conn:
        dbmod.init_db(conn)
        rows = dbmod.list_offers(conn)

    if not rows:
        typer.echo("no offers found.")
        return

    for row in rows:
        who = row["merchant"] or row["external_id"] or "?"
        when = f" ? expires {row['expiry_raw']}" if row["expiry_raw"] else ""
        tag = " [ocr ? verify]" if row["source"] == "ocr" else ""
        typer.echo(f"[{row['id']}] {who}: {row['promo_code']}{when}{tag}")

@txns_app.command("list")
def txns_list(
    bank: str = typer.Argument(
        ..., help="Which bank to list (e.g. hdfc, axis) ? or 'all' for the cross-bank view."
    ),
    db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data dir)."),
    limit: int = typer.Option(50, help="Max number of transactions to show."),
) -> None:
    """List parsed transactions for one bank, or across all banks via the serving view."""
    from muneem import db as dbmod

    db_path = db or load_config().db_path
    if not Path(db_path).exists():
        typer.echo("database not created yet.")
        return

    bank = bank.lower()
    if bank == "all":
        with _ledger(db_path) as conn:
            dbmod.init_db(conn)
            rows = dbmod.list_all_transactions(conn, limit=limit)
        if not rows:
            typer.echo("no transactions found.")
            return
        for row in rows:
            amount = f"+{row['credit']}" if row["credit"] else f"-{row['debit']}"
            part = (row["particulars"] or "")[:40]
            verdict = row["confidence"] or "?"
            typer.echo(
                f"[{row['bank']}] [{row['account_number']}] [{row['txn_date']}]: "
                f"[{part[:40]}] | {amount[:10]} [{verdict}]"
            )
    else:
        with _ledger(db_path) as conn:
            dbmod.init_db(conn)
            rows = dbmod.list_transactions(conn, bank=bank, limit=limit)
        if not rows:
            typer.echo("no transactions found.")
            return
        for row in rows:
            amount = f"+{row['credit']}" if row["credit"] else f"-{row['debit']}"
            part = (row["particulars"] or "")[:40]
            verdict = row["confidence"] or "?"
            typer.echo(
                f"[{row['bank']}] [{row['account_number']}] [{row['txn_date']}]: "
                f"[{part[:40]}] | {amount[:10]} [{verdict}]"
            )

```

```

        return

    if bank not in BANK_DISPATCH:
        typer.echo(f"Unknown bank: {bank}", err=True)
        return

    list_fn, neg_key, pos_key = BANK_DISPATCH[bank]
    with _ledger(db_path) as conn:
        dbmod.init_db(conn)
        rows = list_fn(conn, limit=limit)

    if not rows:
        typer.echo(f"no {bank} transactions found.")
        return

    for row in rows:
        if bank == "axis_cc":
            sign = "+" if row["type"] == "Cr" else "-"
            amount = f"{sign}{row['amount']}"
        else:
            amount = f"+{row[pos_key]}" if row[pos_key] else f"-{row[neg_key]}"
        part = (row["particulars"] or "")[:50]
        # Anything but a clean 'reconciled' verdict is surfaced, never silently listed.
        status = row["doc_status"]
        tag = "" if status == "reconciled" else f"[{status}]"
        typer.echo(f"[{row['account_number']}] {row['txn_date']} : {part:<50} |
{amount:>10}{tag}")

@db_app.command("status")
def db_status(
    db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data dir)."),
) -> None:
    """Show the ledger's encryption state (no statement data is read or printed)."""
    from muneem import db as dbmod
    from muneem.db_key import peek_key, peek_pending_key

    db_path = db or load_config().db_path
    key_state = "present in keychain" if peek_key() else "not yet minted"
    if not Path(db_path).exists():
        typer.echo(f"ledger : {db_path} (does not exist yet)")
        typer.echo(f"key      : {key_state}")
        return
    if dbmod.is_plaintext_sqlite(db_path):
        state = "PLAINTEXT ? run 'muneem db encrypt' to harden it"
    else:
        state = "encrypted (SQLCipher)"
    typer.echo(f"ledger : {db_path}")
    typer.echo(f"state  : {state}")
    typer.echo(f"key    : {key_state}")
    if peek_pending_key():
        typer.echo(
            "pending: a staged rotation key is present (an interrupted 'db rotate-key') ? "
            "the next ledger open reconciles it automatically."
        )

@db_app.command("encrypt")
def db_encrypt(
    db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data dir)."),
) -> None:
    """Migrate an existing *plaintext* ledger to an encrypted one (SQLCipher), in place."""
    from muneem import db as dbmod
    from muneem.db_key import get_or_create_key

    db_path = db or load_config().db_path

```

```

if not Path(db_path).exists():
    typer.echo(
        f"nothing to encrypt: {db_path} does not exist (a new ledger is born encrypted)."
    )
    return
if not dbmod.is_plaintext_sqlite(db_path):
    typer.echo(
        f"{Path(db_path).name} is already encrypted (or not a SQLite file); nothing to do."
    )
    return
n = dbmod.encrypt_database(db_path, get_or_create_key())
typer.echo(
    f"encrypted {Path(db_path).name}: {n} table(s) migrated. A plaintext backup remains at "
    f"{Path(db_path).name}.plaintext.bak ? verify, then securely delete it."
)

def _require_encrypted_ledger(db_path: Path) -> str:
    """Return the ledger's keychain key, or exit with a clear message if it can't be used."""
    from muneem import db as dbmod
    from muneem.db_key import peek_key

    if not Path(db_path).exists() or dbmod.is_plaintext_sqlite(db_path):
        typer.echo("no encrypted ledger here (run 'muneem db encrypt' first).", err=True)
        raise typer.Exit(1)
    key = peek_key()
    if not key:
        typer.echo("no encryption key in the keychain for this ledger.", err=True)
        raise typer.Exit(1)
    return key

@db_app.command("rotate-key")
def db_rotate_key(
    db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data dir)."),
) -> None:
    """Rotate the ledger's encryption key (SQLCipher rekey) and update the keychain."""
    from muneem import db as dbmod
    from muneem import db_key
    from muneem.errors import KeychainError

    db_path = db or load_config().db_path
    current = _require_encrypted_ledger(db_path)
    fresh = db_key.new_key()
    # Both keys keep a durable copy at every instant: stage the fresh key under the PENDING
    # keychain entry, rekey the file, and only then promote it to the primary entry. Any
    # interruption leaves a state that connect_encrypted can heal (file rekeyed -> pending key
    # promoted on next open; file untouched -> primary entry still matches it).
    db_key.set_pending_key(fresh)
    try:
        dbmod.rekey_database(db_path, current, fresh)
    except Exception as exc:
        # Don't trust the exception to tell us how far the rekey got ? check the file.
        if dbmod.key_opens_database(db_path, fresh):
            pass # the rekey actually completed; fall through to promotion below
        elif dbmod.key_opens_database(db_path, current):
            db_key.clear_pending_key()
            typer.echo(
                f"key rotation failed; the ledger is unchanged (still under the existing key): "
                f"{exc}",
                err=True,
            )
            raise typer.Exit(1) from exc
        else:
            typer.echo(

```

```

        "key rotation failed and the ledger opens with neither key. Both keys were kept
"
        "in the keychain ('db-key' = old, 'db-key-pending' = new); restore the ledger "
        "from a backup. Do NOT delete either keychain entry until the ledger opens.",
        err=True,
    )
    raise typer.Exit(1) from exc
try:
    db_key.set_key(fresh) # promote: the file is verifiably under the fresh key now
except KeychainError as exc:
    typer.echo(
        "the ledger was rekeyed but promoting the new key in the keychain failed; the new "
        "key is safe under 'db-key-pending' and the next ledger open will finish the "
        "promotion automatically.",
        err=True,
    )
    raise typer.Exit(1) from exc
db_key.clear_pending_key()
typer.echo(f"rotated the encryption key for {Path(db_path).name}.")
typer.echo(
    "note: backups and any recovery file created BEFORE this rotation still use the old "
    "key. Re-run 'muneem db backup' and 'muneem db setup-recovery' now."
)

@db_app.command("backup")
def db_backup(
    dest: Path = typer.Argument(..., help="Destination path for the encrypted backup copy."),
    db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data dir)."),
) -> None:
    """Write an encrypted online-backup copy of the ledger (consistent even with WAL active)."""
    from muneem import db as dbmod

    db_path = db or load_config().db_path
    key = _require_encrypted_ledger(db_path)
    try:
        dbmod.backup_database(db_path, dest, key)
    except dbmod.EncryptedDatabaseError as exc:
        typer.echo(str(exc), err=True)
        raise typer.Exit(1) from exc
    typer.echo(f"backed up {Path(db_path).name} to {dest} (encrypted).")

@db_app.command("setup-recovery")
def db_setup_recovery(
    out: Path | None = typer.Option(
        None, help="Recovery file path (default: <data dir>/muneem-recovery.json)."
    ),
) -> None:
    """Create a passphrase-wrapped recovery file for the ledger's encryption key.

    Lets you restore the key (and thus the ledger + backups) if the OS keychain is ever lost.
    The
    file holds only the *wrapped* key ? useless without the passphrase.
    """
    import getpass

    from muneem import recovery
    from muneem.db_key import peek_key

    key = peek_key()
    if not key:
        typer.echo("no encryption key in the keychain yet (create the ledger first).", err=True)
        raise typer.Exit(1)
    out_path = out or (load_config().data_dir / "muneem-recovery.json")

```

```

passphrase = getpass.getpass("New recovery passphrase: ")
if passphrase != getpass.getpass("Confirm passphrase: "):
    typer.echo("passphrases did not match.", err=True)
    raise typer.Exit(1)
try:
    recovery.write_recovery_file(out_path, key, passphrase)
except recovery.RecoveryError as exc:
    typer.echo(str(exc), err=True)
    raise typer.Exit(1) from exc
typer.echo(
    f"wrote recovery file to {out_path}. Store it OFF this machine (password manager / "
    "separate drive) and remember the passphrase ? you need BOTH to recover. Without them, a
"
    "lost keychain means the ledger and its backups are unrecoverable."
)

```

```

@db_app.command("recover")
def db_recover(
    file: Path = typer.Argument(..., exists=True, dir_okay=False, help="Recovery file path."),
    db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data dir)."),
) -> None:
    """Restore the ledger's encryption key from a recovery file + passphrase (into the
    keychain)."""
    import getpass

    from muneem import db as dbmod
    from muneem import db_key, recovery

    passphrase = getpass.getpass("Recovery passphrase: ")
    try:
        key = recovery.read_recovery_file(file, passphrase)
    except recovery.RecoveryError as exc:
        typer.echo(str(exc), err=True)
        raise typer.Exit(1) from exc
    db_path = db or load_config().db_path
    # If the ledger is present, verify the recovered key actually opens it before committing it.
    if Path(db_path).exists() and not dbmod.is_plaintext_sqlite(db_path):
        try:
            with dbmod.connect(db_path, key=key) as conn:
                conn.execute("SELECT 1")
        except dbmod.EncryptedDatabaseError as exc:
            typer.echo("recovered a key, but it does not open the current ledger.", err=True)
            raise typer.Exit(1) from exc
    db_key.set_key(key)
    typer.echo("recovered the encryption key into the keychain.")

if __name__ == "__main__":
    app()

```

## 12. assistant (2026-06-25T16:42:41.944Z)

Now let me read all the new split modules to verify behavior is preserved.

## 13. user (2026-06-25T16:42:42.892Z)

```

1     """The ``muneem`` command-line interface ? re-exports from the ``cli/`` subpackage.
2
3     See ``muneem/cli/`` for the module layout. The entry point (``pyproject.toml``
4     ``[project.scripts]``) still points to ``muneem.cli:app``.
5     """
6
7     from muneem.cli import app # noqa: F401
8

```

#### 14. user (2026-06-25T16:42:43.362Z)

```
1      """The ``muneem`` command-line interface.
2
3      Subcommands are split across focused modules: ``ingest``, ``offers``, ``txns``,
4      ``eval``, and ``db`` ? each in its own file under ``cli/``. This module assembles
5      the Typer app.
6
7      ``fetch`` (Gmail pull) lands with Phase F.
8      """
9
10     from __future__ import annotations
11
12     import typer
13
14     from muneem import __version__
15     from muneem.cli.shared import _version_callback
16     from muneem.config import load_config
17
18     app = typer.Typer(
19         add_completion=False,
20         help="muneem ? a local-first personal-finance bookkeeper.",
21         no_args_is_help=True,
22     )
23
24     # Sub-apps ? each wired in the module that defines its commands.
25     from muneem.cli.db import db_app # noqa: E402
26     from muneem.cli.offers import offers_app # noqa: E402
27     from muneem.cli.txns import txns_app # noqa: E402
28
29     app.add_typer(offers_app, name="offers")
30     app.add_typer(txns_app, name="txns")
31     app.add_typer(db_app, name="db")
32
33
34     @app.callback()
35     def main(
36         version: bool = typer.Option(
37             False,
38             "--version",
39             callback=_version_callback,
40             is_eager=True,
41             help="Show the muneem version and exit.",
42         ),
43     ) -> None:
44         """muneem command-line interface."""
45
46
47     @app.command()
48     def info() -> None:
49         """Show resolved local paths (no secrets, no sensitive data)."""
50         cfg = load_config()
51         present = "present" if cfg.config_file.is_file() else "absent"
52         typer.echo(f"muneem {__version__}")
53         typer.echo(f"config dir : {cfg.config_dir}")
54         typer.echo(f"data dir   : {cfg.data_dir}")
55         typer.echo(f"db path   : {cfg.db_path}")
56         typer.echo(f"staging   : {cfg.staging_dir}")
57         typer.echo(f"config file: {cfg.config_file} ({present})")
58
59
60     # Command registrations ? registered explicitly to avoid circular decorator imports.
61     from muneem.cli.eval import eval_statements # noqa: E402
62     from muneem.cli.ingest import ingest # noqa: E402
63
```

```

64     app.command()(ingest)
65     app.command("eval")(eval_statements)
66     # side-effects at import time. E402 is suppressed because the imports must follow
67     # the app/sub-app definitions above.
68     from muneem.cli.eval import eval_statements # noqa: E402, F401
69     from muneem.cli.ingest import ingest # noqa: E402, F401
70

```

#### 15. user (2026-06-25T16:42:43.522Z)

```

1     """Shared CLI utilities ? no imports from ``muneem.cli`` to avoid circularity."""
2
3     from __future__ import annotations
4
5     from contextlib import contextmanager
6     from pathlib import Path
7
8     import typer
9
10    from muneem import __version__
11    from muneem import db as dbmod
12
13
14    @contextmanager
15    def _ledger(db_path: Path):
16        """Open the encrypted ledger, turning an un-migrated plaintext file into a clean CLI
17        error."""
18        try:
19            with dbmod.connect_encrypted(db_path) as conn:
20                yield conn
21        except dbmod.EncryptedDatabaseError as exc:
22            typer.echo(str(exc), err=True)
23            raise typer.Exit(1) from exc
24
25    # Dispatch table for supported banks (extensible as new parsers land).
26    BANK_DISPATCH = {
27        "axis": (dbmod.list_axis_transactions, "debit", "credit"),
28        "axis_cc": (dbmod.list_axis_cc_transactions, "amount", "amount"),
29        "au": (dbmod.list_au_transactions, "debit", "credit"),
30        "hdfc": (dbmod.list_hdfc_transactions, "withdrawal", "deposit"),
31    }
32
33
34    def _version_callback(value: bool) -> None:
35        if value:
36            typer.echo(f"muneem {__version__}")
37            raise typer.Exit()
38

```

#### 16. assistant (2026-06-25T16:42:58.750Z)

Let me read the remaining split modules.

#### 17. user (2026-06-25T16:42:59.242Z)

```

1     """The ``ingest`` command ? parse a PDF into the local store."""
2
3     from __future__ import annotations
4
5     from pathlib import Path
6
7     import typer
8

```



```

9      from muneem.cli._shared import _ledger
10     from muneem.config import load_config
11     from muneem.log import get_logger
12
13     _log = get_logger(__name__)
14
15
16     def _ingest_bank_statement(pdf: Path, db_path: Path, doc_type: str, *, interactive:
bool) -> None:
17         """Unlock (if encrypted) ? extract ? select parser ? store reconciled
transactions."""
18         from muneem import db as dbmod
19         from muneem.extract import extract
20         from muneem.parsers import get_parser
21         from muneem.statement_unlock import resolve_password
22         from muneem.unlock import is_encrypted
23
24         password: str | None = ""
25         if is_encrypted(pdf):
26             password = resolve_password(pdf, interactive=interactive)
27             if password is None:
28                 hint = " (run without --no-prompt to enter it)" if not interactive else ""
29                 typer.echo(f"Could not unlock {pdf.name}: no working password{hint}.",
err=True)
30                 raise typer.Exit(1)
31
32         ex = extract(pdf, password=password or "")
33         parser = get_parser(ex.text)
34         if parser is None:
35             typer.echo(f"No parser recognizes {pdf.name} as a supported bank statement.",
err=True)
36             raise typer.Exit(1)
37
38         with _ledger(db_path) as conn:
39             dbmod.init_db(conn)
40             existing = dbmod.get_document_by_sha(conn, ex.sha256)
41             if existing is not None and existing["status"] == "reconciled":
42                 typer.echo(f"already ingested (doc {existing['id']}): {pdf.name}")
43                 return
44             if existing is not None:
45                 doc_id = int(existing["id"])
46                 dbmod.clear_document_rows(conn, doc_id)
47                 typer.echo(
48                     f"re-parsing previously {existing['status']} document (doc {doc_id}):
{pdf.name}"
49                 )
50             else:
51                 doc_id = dbmod.insert_document(
52                     conn,
53                     source="file",
54                     external_id=pdf.name,
55                     doc_type=doc_type,
56                     sha256=ex.sha256,
57                     path=str(pdf.resolve()),
58                     status="parsed",
59                 )
60                 count = parser.parse_and_insert(pdf, conn, doc_id, password=password)
61                 doc = dbmod.get_document_by_sha(conn, ex.sha256)
62                 status = doc["status"] if doc else "parsed"
63
64             typer.echo(
65                 f"ingested {pdf.name} (doc {doc_id}); parsed {count} transaction(s) "
66                 f"via {parser.__class__.__name__} [{status}]"
67             )
68

```

```

69
70 def ingest(
71     pdf: Path = typer.Argument(
72         ..., exists=True, dir_okay=False, readable=True, help="Path to a PDF to ingest."
73     ),
74     doc_type: str = typer.Option("promo", help="Document category."),
75     db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data
dir)."),
76     ocr: bool = typer.Option(
77         False, "--ocr", help="Also OCR images locally (Tesseract) for codes in banners."
78     ),
79     no_prompt: bool = typer.Option(
80         False,
81         "--no-prompt",
82         help="Never prompt for a statement password; use only cached/configured
candidates.",
83     ),
84 ) -> None:
85     """Store a PDF's offers (promo) or transactions (bank_statement) in the local DB."""
86     from muneem import db as dbmod
87     from muneem.extract import extract
88     from muneem.offers import extract_offer
89
90     db_path = db or load_config().db_path
91
92     if doc_type == "bank_statement":
93         _ingest_bank_statement(pdf, db_path, doc_type, interactive=not no_prompt)
94         return
95
96     ex = extract(pdf)
97     text_fields = extract_offer(ex.text)
98     found: list[tuple[str, str]] = [(c, "text") for c in text_fields.codes]
99     expiry_raw, expiry_type = text_fields.expiry_raw, text_fields.expiry_type
100
101     if ocr:
102         from muneem.ocr import ocr_pdf, tesseract_available
103
104         if not tesseract_available():
105             typer.echo("warning: --ocr requested but Tesseract not found; skipping
OCR.", err=True)
106         else:
107             ocr_fields = extract_offer(ocr_pdf(pdf))
108             known = {c for c, _ in found}
109             for code in ocr_fields.codes:
110                 if code not in known:
111                     found.append((code, "ocr"))
112                     known.add(code)
113             if expiry_raw is None:
114                 expiry_raw, expiry_type = ocr_fields.expiry_raw, ocr_fields.expiry_type
115
116     with _ledger(db_path) as conn:
117         dbmod.init_db(conn)
118         existing = dbmod.get_document_by_sha(conn, ex.sha256)
119         if existing is not None:
120             typer.echo(f"already ingested (doc {existing['id']}): {pdf.name}")
121             return
122         doc_id = dbmod.insert_document(
123             conn,
124             source="file",
125             external_id=pdf.name,
126             doc_type=doc_type,
127             sha256=ex.sha256,
128             path=str(pdf.resolve()),
129             status="parsed",
130         )

```

```

131         for code, src in found:
132             dbmod.insert_offer(
133                 conn,
134                 document_id=doc_id,
135                 promo_code=code,
136                 expiry_raw=expiry_raw,
137                 expiry_type=expiry_type,
138                 source=src,
139                 source_text=ex.text[:2000],
140             )
141
142     summary = ", ".join(f"{c} [{s}]" for c, s in found) if found else "none"
143     typer.echo(f"ingested {pdf.name} (doc {doc_id}); {len(found)} offer(s): {summary}")
144

```

## 18. user (2026-06-25T16:42:59.619Z)

```

1     """The ``eval`` command ? evaluate parsers against a folder of statements."""
2
3     from __future__ import annotations
4
5     from pathlib import Path
6
7     import typer
8
9     from muneem import db as dbmod
10    from muneem.extract import extract
11    from muneem.log import get_logger
12    from muneem.parsers import get_parser
13    from muneem.redaction import redact_filename
14    from muneem.statement_unlock import resolve_password
15    from muneem.unlock import is_encrypted
16
17    _log = get_logger(__name__)
18
19
20    def _eval_statement(pdf: Path, *, interactive: bool) -> tuple[str, int, str | None]:
21        """Run one statement through the full parse pipeline into an in-memory DB.
22
23        Returns ``(outcome, row_count, parser_name)`` where ``outcome`` is the document
24        status
25        (`reconciled` / `unreconciled`) or a failure tag (`locked` / `no-parser` /
26        `error`).
27
28        """
29        password = ""
30        if is_encrypted(pdf):
31            password = resolve_password(pdf, interactive=interactive) or ""
32            if not password:
33                return ("locked", 0, None)
34
35        try:
36            ex = extract(pdf, password=password)
37            parser = get_parser(ex.text)
38        except Exception as exc:
39            _log.warning(
40                "eval: extract/route failed for %s (%s)", redact_filename(pdf),
41                type(exc).__name__
42            )
43            return ("error", 0, None)
44
45        if parser is None:
46            return ("no-parser", 0, None)
47
48        name = parser.__class__.__name__
49
50        try:
51            with dbmod.connect(":memory:") as conn:
52                dbmod.init_db(conn)
53                doc_id = dbmod.insert_document(

```

```

46         conn, source="file", external_id=pdf.name, sha256=ex.sha256,
status="parsed"
47     )
48     count = parser.parse_and_insert(pdf, conn, doc_id, password=password)
49     doc = dbmod.get_document_by_sha(conn, ex.sha256)
50     status = doc["status"] if doc else "parsed"
51 except Exception as exc:
52     _log.warning(
53         "eval: parse/store failed for %s (%s)", redact_filename(pdf),
type(exc).__name__
54     )
55     return ("error", 0, name)
56     return (status, count, name)
57
58
59 def eval_statements(
60     folder: Path = typer.Argument(
61         ...,
62         exists=True,
63         file_okay=False,
64         readable=True,
65         help="A LOCAL folder of statement PDFs to evaluate.",
66     ),
67     prompt: bool = typer.Option(
68         False,
69         "--prompt/--no-prompt",
70         help="Prompt for a password when cached/configured candidates don't unlock a
file.",
71     ),
72 ) -> None:
73     """Evaluate the parsers against a folder of statements and report the reconcile
rate."""
74     pdfs = sorted(folder.glob("*.pdf"))
75     if not pdfs:
76         typer.echo(f"no .pdf files in {folder}")
77         return
78     reconciled = 0
79     for pdf in pdfs:
80         outcome, count, name = _eval_statement(pdf, interactive=prompt)
81         reconciled += outcome == "reconciled"
82         typer.echo(f" {outcome:12} {count:4d} row(s) {name or '?':20} {pdf.name}")
83     typer.echo(f"\n{reconciled}/{len(pdfs)} reconciled")
84

```

## 19. user (2026-06-25T16:43:16.797Z)

```

1     """The ``db`` subcommand group ? ledger encryption lifecycle."""
2
3     from __future__ import annotations
4
5     from pathlib import Path
6
7     import typer
8
9     from muneem import db as dbmod
10    from muneem.config import load_config
11
12    db_app = typer.Typer(no_args_is_help=True, help="Manage the local encrypted ledger.")
13
14
15    def _require_encrypted_ledger(db_path: Path) -> str:
16        """Return the ledger's keychain key, or exit with a clear message if it can't be
used."""
17        from muneem.db_key import peek_key
18

```

```

19         if not Path(db_path).exists() or dbmod.is_plaintext_sqlite(db_path):
20             typer.echo("no encrypted ledger here (run 'muneem db encrypt' first).",
err=True)
21             raise typer.Exit(1)
22         key = peek_key()
23         if not key:
24             typer.echo("no encryption key in the keychain for this ledger.", err=True)
25             raise typer.Exit(1)
26         return key
27
28
29     @db_app.command("status")
30     def db_status(
31         db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data
dir)."),
32     ) -> None:
33         """Show the ledger's encryption state (no statement data is read or printed)."""
34         from muneem.db_key import peek_key, peek_pending_key
35
36         db_path = db or load_config().db_path
37         key_state = "present in keychain" if peek_key() else "not yet minted"
38         if not Path(db_path).exists():
39             typer.echo(f"ledger : {db_path} (does not exist yet)")
40             typer.echo(f"key      : {key_state}")
41             return
42         if dbmod.is_plaintext_sqlite(db_path):
43             state = "PLAINTEXT ? run 'muneem db encrypt' to harden it"
44         else:
45             state = "encrypted (SQLCipher)"
46             typer.echo(f"ledger : {db_path}")
47             typer.echo(f"state  : {state}")
48             typer.echo(f"key    : {key_state}")
49         if peek_pending_key():
50             typer.echo(
51                 "pending: a staged rotation key is present (an interrupted 'db rotate-key'?)
? "
52                 "the next ledger open reconciles it automatically."
53             )
54
55
56     @db_app.command("encrypt")
57     def db_encrypt(
58         db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data
dir)."),
59     ) -> None:
60         """Migrate an existing *plaintext* ledger to an encrypted one (SQLCipher), in
place."""
61         from muneem.db_key import get_or_create_key
62
63         db_path = db or load_config().db_path
64         if not Path(db_path).exists():
65             typer.echo(
66                 f"nothing to encrypt: {db_path} does not exist (a new ledger is born
encrypted)."
67             )
68             return
69         if not dbmod.is_plaintext_sqlite(db_path):
70             typer.echo(
71                 f"{Path(db_path).name} is already encrypted (or not a SQLite file); nothing
to do."
72             )
73             return
74         n = dbmod.encrypt_database(db_path, get_or_create_key())
75         typer.echo(
76             f"encrypted {Path(db_path).name}: {n} table(s) migrated. A plaintext backup

```

```

remains at "
77         f"{Path(db_path).name}.plaintext.bak ? verify, then securely delete it."
78     )
79
80
81     @db_app.command("rotate-key")
82     def db_rotate_key(
83         db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data
dir)."),
84     ) -> None:
85         """Rotate the ledger's encryption key (SQLCIPHER rekey) and update the keychain."""
86         from muneem import db_key
87         from muneem.errors import KeychainError
88
89         db_path = db or load_config().db_path
90         current = _require_encrypted_ledger(db_path)
91         fresh = db_key.new_key()
92         db_key.set_pending_key(fresh)
93         try:
94             dbmod.rekey_database(db_path, current, fresh)
95         except Exception as exc:
96             if dbmod.key_opens_database(db_path, fresh):
97                 pass
98             elif dbmod.key_opens_database(db_path, current):
99                 db_key.clear_pending_key()
100                 typer.echo(
101                     f"key rotation failed; the ledger is unchanged (still under the existing
key): "
102                     f"{exc}",
103                     err=True,
104                 )
105                 raise typer.Exit(1) from exc
106             else:
107                 typer.echo(
108                     "key rotation failed and the ledger opens with neither key. Both keys
were kept "
109                     "in the keychain ('db-key' = old, 'db-key-pending' = new); restore the
ledger "
110                     "from a backup. Do NOT delete either keychain entry until the ledger
opens.",
111                     err=True,
112                 )
113                 raise typer.Exit(1) from exc
114         try:
115             db_key.set_key(fresh)
116         except KeychainError as exc:
117             typer.echo(
118                 "the ledger was rekeyed but promoting the new key in the keychain failed;
the new "
119                 "key is safe under 'db-key-pending' and the next ledger open will finish the
"
120                 "promotion automatically.",
121                 err=True,
122             )
123             raise typer.Exit(1) from exc
124         db_key.clear_pending_key()
125         typer.echo(f"rotated the encryption key for {Path(db_path).name}.")
126         typer.echo(
127             "note: backups and any recovery file created BEFORE this rotation still use the
old "
128             "key. Re-run 'muneem db backup' and 'muneem db setup-recovery' now."
129         )
130
131
132     @db_app.command("backup")

```

```

133     def db_backup(
134         dest: Path = typer.Argument(..., help="Destination path for the encrypted backup
copy."),
135         db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data
dir)."),
136     ) -> None:
137         """Write an encrypted online-backup copy of the ledger (consistent even with WAL
active)."""
138         db_path = db or load_config().db_path
139         key = _require_encrypted_ledger(db_path)
140         try:
141             dbmod.backup_database(db_path, dest, key)
142         except dbmod.EncryptedDatabaseError as exc:
143             typer.echo(str(exc), err=True)
144             raise typer.Exit(1) from exc
145         typer.echo(f"backed up {Path(db_path).name} to {dest} (encrypted).")
146
147
148     @db_app.command("setup-recovery")
149     def db_setup_recovery(
150         out: Path | None = typer.Option(
151             None, help="Recovery file path (default: <data dir>/muneem-recovery.json)."
152         ),
153     ) -> None:
154         """Create a passphrase-wrapped recovery file for the ledger's encryption key."""
155         import getpass
156
157         from muneem import recovery
158         from muneem.db_key import peek_key
159
160         key = peek_key()
161         if not key:
162             typer.echo("no encryption key in the keychain yet (create the ledger first).",
err=True)
163             raise typer.Exit(1)
164         out_path = out or (load_config().data_dir / "muneem-recovery.json")
165         passphrase = getpass.getpass("New recovery passphrase: ")
166         if passphrase != getpass.getpass("Confirm passphrase: "):
167             typer.echo("passphrases did not match.", err=True)
168             raise typer.Exit(1)
169         try:
170             recovery.write_recovery_file(out_path, key, passphrase)
171         except recovery.RecoveryError as exc:
172             typer.echo(str(exc), err=True)
173             raise typer.Exit(1) from exc
174         typer.echo(
175             f"wrote recovery file to {out_path}. Store it OFF this machine (password manager
/ "
176             "separate drive) and remember the passphrase ? you need BOTH to recover. Without
them, "
177             "a lost keychain means the ledger and its backups are unrecoverable."
178         )
179
180
181     @db_app.command("recover")
182     def db_recover(
183         file: Path = typer.Argument(..., exists=True, dir_okay=False, help="Recovery file
path."),
184         db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data
dir)."),
185     ) -> None:
186         """Restore the ledger's encryption key from a recovery file + passphrase (into the
keychain)."""
187         import getpass
188

```

```

189     from muneem import db_key, recovery
190
191     passphrase = getpass.getpass("Recovery passphrase: ")
192     try:
193         key = recovery.read_recovery_file(file, passphrase)
194     except recovery.RecoveryError as exc:
195         typer.echo(str(exc), err=True)
196         raise typer.Exit(1) from exc
197     db_path = db or load_config().db_path
198     if Path(db_path).exists() and not dbmod.is_plaintext_sqlite(db_path):
199         try:
200             with dbmod.connect(db_path, key=key) as conn:
201                 conn.execute("SELECT 1")
202             except dbmod.EncryptedDatabaseError as exc:
203                 typer.echo("recovered a key, but it does not open the current ledger.",
err=True)
204                 raise typer.Exit(1) from exc
205         db_key.set_key(key)
206         typer.echo("recovered the encryption key into the keychain.")
207

```

## 20. user (2026-06-25T16:43:17.262Z)

```

1     """The ``offers`` subcommand group."""
2
3     from __future__ import annotations
4
5     from pathlib import Path
6
7     import typer
8
9     from muneem import db as dbmod
10    from muneem.cli._shared import _ledger
11    from muneem.config import load_config
12
13    offers_app = typer.Typer(no_args_is_help=True, help="Work with extracted offers.")
14
15
16    @offers_app.command("list")
17    def offers_list(
18        db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data
dir)."),
19    ) -> None:
20        """List extracted offers."""
21        db_path = db or load_config().db_path
22        if not Path(db_path).exists():
23            typer.echo("no offers yet (database not created).")
24            return
25
26        with _ledger(db_path) as conn:
27            dbmod.init_db(conn)
28            rows = dbmod.list_offers(conn)
29
30        if not rows:
31            typer.echo("no offers found.")
32            return
33        for row in rows:
34            who = row["merchant"] or row["external_id"] or "?"
35            when = f" ? expires {row['expiry_raw']}" if row["expiry_raw"] else ""
36            tag = " [ocr ? verify]" if row["source"] == "ocr" else ""
37            typer.echo(f"[{row['id']}] {who}: {row['promo_code']}{when}{tag}")
38

```

## 21. user (2026-06-25T16:43:17.435Z)



```

1      """The ``txns`` subcommand group."""
2
3      from __future__ import annotations
4
5      from pathlib import Path
6
7      import typer
8
9      from muneem import db as dbmod
10     from muneem.cli._shared import BANK_DISPATCH, _ledger
11     from muneem.config import load_config
12
13     txns_app = typer.Typer(no_args_is_help=True, help="Work with parsed transactions.")
14
15
16     @txns_app.command("list")
17     def txns_list(
18         bank: str = typer.Argument(
19             ..., help="Which bank to list (e.g. hdfc, axis) ? or 'all' for the cross-bank
view."
20         ),
21         db: Path | None = typer.Option(None, help="SQLite path (defaults to the local data
dir)."),
22         limit: int = typer.Option(50, help="Max number of transactions to show."),
23     ) -> None:
24         """List parsed transactions for one bank, or across all banks via the serving
view."""
25         db_path = db or load_config().db_path
26         if not Path(db_path).exists():
27             typer.echo("database not created yet.")
28             return
29
30         bank = bank.lower()
31         if bank == "all":
32             with _ledger(db_path) as conn:
33                 dbmod.init_db(conn)
34                 rows = dbmod.list_all_transactions(conn, limit=limit)
35             if not rows:
36                 typer.echo("no transactions found.")
37                 return
38             for row in rows:
39                 amount = f"+{row['credit']}" if row["credit"] else f"-{row['debit']}"
40                 part = (row["particulars"] or "")[:40]
41                 verdict = row["confidence"] or "?"
42                 typer.echo(
43                     f"[{row['bank']}] [{row['account_number']}] [{row['txn_date']}] : "
44                     f"[{part:<40}] | [{amount:>10}] [{verdict}]"
45                 )
46             return
47
48         if bank not in BANK_DISPATCH:
49             typer.echo(f"Unknown bank: {bank}", err=True)
50             return
51
52         list_fn, neg_key, pos_key = BANK_DISPATCH[bank]
53         with _ledger(db_path) as conn:
54             dbmod.init_db(conn)
55             rows = list_fn(conn, limit=limit)
56
57         if not rows:
58             typer.echo(f"no {bank} transactions found.")
59             return
60         for row in rows:
61             if bank == "axis_cc":
62                 sign = "+" if row["type"] == "Cr" else "-"

```

```

63         amount = f"{{sign}}{{row['amount']}}"
64     else:
65         amount = f"+{{row[pos_key]}}" if row[pos_key] else f"-{{row[neg_key]}}"
66     part = (row["particulars"] or "")[:50]
67     status = row["doc_status"]
68     tag = "" if status == "reconciled" else f" [{{status}}]"
69     typer.echo(f"{{row['account_number']}} {{row['txn_date']}} : {{part:<50}} |
{{amount:>10}}{{tag}}")
70

```

## 22. assistant (2026-06-25T16:43:30.655Z)

All the `db`/`offers`/`txns` modules are faithful to the original. I found two things worth checking. Let me verify how the tests monkeypatch `extract` ? because `eval.py` hoisted those imports to module level while `ingest.py` kept them lazy, and the PR explicitly claims lazy imports were preserved for monkeypatching.

## 23. user (2026-06-25T16:43:54.288Z)

```

=== monkeypatch of extract / get_parser in tests ===
tests/test_cli.py:24: def test_ingest_bank_statement_unlocks_and_passes_password(tmp_path,
monkeypatch):
tests/test_cli.py:30:     monkeypatch.setattr("muneem.unlock.is_encrypted", lambda p: True)
tests/test_cli.py:31:     monkeypatch.setattr(
tests/test_cli.py:39:     monkeypatch.setattr("muneem.extract.extract", fake_extract)
tests/test_cli.py:46:     monkeypatch.setattr("muneem.parsers.get_parser", lambda text:
FakeParser())
tests/test_cli.py:56: def test_quarantined_document_can_be_reingested_after_parser_fix(tmp_path,
monkeypatch):
tests/test_cli.py:65:     monkeypatch.setattr("muneem.unlock.is_encrypted", lambda p: False)
tests/test_cli.py:66:     monkeypatch.setattr(
tests/test_cli.py:100:     monkeypatch.setattr("muneem.parsers.get_parser", lambda text:
BrokenParser())
tests/test_cli.py:105:     monkeypatch.setattr("muneem.parsers.get_parser", lambda text:
FixedParser())
tests/test_cli.py:150: def test_ingest_bank_statement_unlock_failure_exits_nonzero(tmp_path,
monkeypatch):
tests/test_cli.py:153:     monkeypatch.setattr("muneem.unlock.is_encrypted", lambda p: True)
tests/test_cli.py:154:     monkeypatch.setattr(
Binary file tests/__pycache__/test_cli.cpython-313-pytest-9.0.3.pyc matches

```

## 24. user (2026-06-25T16:43:54.288Z)

```

=== eval-related tests ===
test_eval_cli.py
--- patch targets in eval tests ---

```

## 25. assistant (2026-06-25T16:44:00.977Z)

The ingest tests patch `muneem.extract.extract` (the source module) ? which works with `ingest.py`'s lazy imports. Let me check how the eval tests patch, since `eval.py` now binds those names at module level.

## 26. user (2026-06-25T16:44:01.614Z)

```

1     """Tests for `muneem eval` ? the seed?golden onboarding harness.
2
3     Runs the CLI against a folder of *synthetic* statements (no secrets, no real data) and
checks the
4     reconcile-rate report. Reuses the born-digital fixtures from the regression suite.
5     """
6
7     from typer.testing import CliRunner
8

```

```

9      from muneem.cli import app
10     from test_regression_synt import _au_savings_pdf, _hdfc_pdf
11
12     runner = CliRunner()
13
14
15     def test_eval_reports_reconcile_rate(tmp_path):
16         _hdfc_pdf(tmp_path / "hdfc.pdf")
17         _au_savings_pdf(tmp_path / "au.pdf")
18         result = runner.invoke(app, ["eval", str(tmp_path)])
19         assert result.exit_code == 0, result.output
20         assert "2/2 reconciled" in result.output
21         assert "reconciled" in result.output # per-file outcome line
22
23
24     def test_eval_empty_folder(tmp_path):
25         result = runner.invoke(app, ["eval", str(tmp_path)])
26         assert result.exit_code == 0
27         assert "no .pdf files" in result.output
28

```

## 27. assistant (2026-06-25T16:44:36.423Z)

The eval test uses real synthetic PDFs and does **not** monkeypatch `extract` at all ? so `eval.py`'s module-level imports are safe for the current suite. Now let me check the docs change and confirm the full command surface is intact.

## 28. user (2026-06-25T16:44:39.198Z)

```
=== docs/code-health diff ===
```

```

=== registered command surface (live introspection) ===
Traceback (most recent call last):
  File "<string>", line 2, in <module>
    from muneem.cli import app
ModuleNotFoundError: No module named 'muneem'

```

## 29. user (2026-06-25T16:44:50.272Z)

```
=== docs diff (two-dot) ===
```

## 30. user (2026-06-25T16:44:53.321Z)

```

using .venv/Scripts/python.exe
info
ingest
eval
[offers]
  list
[txns]
  list
[db]
  status
  encrypt
  rotate-key
  backup
  setup-recovery
  recover

```

## 31. assistant (2026-06-25T16:45:10.436Z)

Command tree is fully intact ? all 12 commands present (only cosmetic change: `ingest`/`eval` registration order swapped vs the original). Let me grab the docs change directly from the PR since the local diff base is muddled by main having moved.

## 32. user (2026-06-25T16:45:12.891Z)

```
diff --git a/docs/code-health-2026-06-25.md b/docs/code-health-2026-06-25.md
index 5d1e236..1ad44c0 100644
--- a/docs/code-health-2026-06-25.md
+++ b/docs/code-health-2026-06-25.md
@@ -205,15 +205,16 @@ This is the project's strongest dimension. Every DESIGN.md security rule
has a c

## 9. Prioritized Recommendations

-### ? High priority (do now, high impact, low effort)
+> **Progress (2026-06-25):** All 4 ? high-priority items + 1 CI fix completed (PRs #74?#79).
+> Now working on ? medium-priority items.

-1. **Refresh `CURRENT_STATE.md`** ? it's the first doc a contributor reads. Update the "Next
Steps" section to reflect actual current state (post #50?#68 hardening). Remove or update "Do
now" items that are already done.
+### ? High priority (do now, high impact, low effort) ? ? ALL DONE

-2. **Fix DESIGN.md § 8.3 contradiction** ? the "Still parked" sentence contradicts the resolved
block directly below it. Remove the stale line.
-
-3. **Move IMAP/Outlook stubs out of the main source tree** ? create
`src/muneem/ingestion/experimental/` and move `imap.py` and `outlook.py` there. Update
`ingestion/__init__.py` to not export them. This keeps `from muneem.ingestion import ...` honest
about what's supported.
-
-4. **Resolve `safe_amount()` vs `parse_amount()`** ? rename `safe_amount` to `safe_flow_amount`
(or `parse_flow_amount`) and add a deprecation path. The current naming is actively misleading
for balance columns. Add a test that verifies `parse_amount("0.00") == 0.0` (should already
exist) and `safe_amount("0.00") is None` (document the divergence explicitly).
+1. ? **Refresh `CURRENT_STATE.md`** ? Done (PR #76). Updated date, marked re-verify + in-
flight-PR-review as done, linked this report.
+2. ? **Fix DESIGN.md § 8.3 contradiction** ? Done (PR #75). "Items 8.3 and 8.5 remain open" ?
"8.3 is resolved; 8.5 remains open."
+3. ? **Move IMAP/Outlook stubs out of the main source tree** ? Done (PR #77). Moved to
`ingestion/experimental/`, removed from public API, narrowed mypy override, added `assert`
guards in `gmail.py`.
+4. ? **Resolve `safe_amount()` vs `parse_amount()`** ? Done (PR #78). Renamed to
`safe_flow_amount()` with deprecation shim. 4 new tests documenting the zero-divergence.
+5. ? ***(Bonus) Fix CI ? OS-agnostic constraints + ruff errors** ? Done (PR #79). Removed `-c
constraints.txt` from CI; fixed pre-existing `E501` errors in `tools/evaluate_odl.py` that were
breaking all PRs.

### ? Medium priority (do soon, moderate effort)
```

```
diff --git a/src/muneem/cli.py b/src/muneem/cli.py
```

## 33. assistant (2026-06-25T16:46:00.201Z)

I've completed the sync and a full review. Here's the summary.

### Sync

```
- **`badminton-highlight-indexer`** (primary): already up to date with `origin/master`; fetched
one new remote branch (`o1-p11/gemini-remote-file-gc`). Working tree clean.
- **`muneem`**: PR branch `refactor/split-cli` is checked out locally and already at
`origin/refactor/split-cli` (`ec091d5`). Up to date.
```

## Review of [muneem#80](https://github.com/avidullu/muneem/pull/80) ? split `cli.py` into a `cli` subpackage

**\*\*Verdict:** a faithful, behavior-preserving split. CI is green across all platforms. One small cleanup I'd want before merge, plus a consistency nit.\*\*

I confirmed the split is semantically faithful by diffing the original `cli.py` (from `origin/main`) against each new module `_shared.py`, `ingest.py`, `eval.py`, `db.py`, `offers.py`, `txns.py` all carry the original command bodies verbatim. Live introspection confirms all 12 commands are still registered (`info`, `ingest`, `eval`, `offers list`, `txns list`, `db {status,encrypt,rotate-key,backup,setup-recovery,recover}`), and the entry point is unchanged (`muneem.cli:app` resolves through the thin re-export). I relied on the **green CI** (Ubuntu 3.11/3.12/3.13 + Windows 3.13, 478 tests, ruff + mypy clean) for the test/lint/type gates rather than re-running them.

## Findings

**\*\*1.** Dead duplicate imports + an orphaned comment ? `src/muneem/cli/__init__.py:66-69` (should fix)\*\*

```
python
60 # Command registrations ? registered explicitly to avoid circular decorator imports.
61 from muneem.cli.eval import eval_statements # noqa: E402
62 from muneem.cli.ingest import ingest        # noqa: E402
63
64 app.command()(ingest)
65 app.command("eval")(eval_statements)
66 # side-effects at import time. E402 is suppressed because the imports must follow
67 # the app/sub-app definitions above.
68 from muneem.cli.eval import eval_statements # noqa: E402, F401 ? dead
69 from muneem.cli.ingest import ingest        # noqa: E402, F401 ? dead
...
```

Lines 68?69 re-import names already imported (61?62) and already used (64?65), so they're pure no-ops marked `# noqa: F401`. The comment on 66?67 reads mid-sentence ("side-effects at import time?") ? its opening was clearly deleted in a botched edit. **\*\*Ruff won't catch this\*\*** (F811 doesn't fire because the first binding `*is*` used before the redefinition, and F401 is suppressed), so it slipped past green CI. Functionally harmless, but it's dead code in a PR whose entire point is a clean split. Fix = delete lines 66?69.

**\*\*2.** `eval.py` breaks the stated "lazy imports" rule ? `src/muneem/cli/eval.py:9-15` (consistency nit)\*\*

The PR description lists **"Lazy imports: `extract()` imported inside functions to preserve test monkeypatching"** as a key design decision, and `ingest.py` honors it (imports stay inside the functions). But `eval.py` hoists `extract`, `get_parser`, `resolve_password`, `is_encrypted` to **module level**. It's safe *today* ? `test_eval_cli.py` runs real synthetic PDFs and never monkeypatches `extract` ? but it's a latent trap: an eval test that later patches `muneem.extract.extract` (the way `test_cli.py:39` does for `ingest`) would silently have no effect. Either make eval lazy for consistency, or leave a note that eval intentionally diverges.

**\*\*3.** Cosmetic:\*\* registration order swapped `info, eval, ingest` ? `info, ingest, eval`, so `--help` lists them in a different order. No functional impact.

**\*\*Docs:\*\*** the `code-health-2026-06-25.md` change is a benign tracker update (flips the 4 ? items to ? with PR refs) ? accurate, non-functional.

## Recommendation

Clean up **\*\*#1\*\*** (4-line deletion; optionally address #2), then it's good to merge ? the split itself is correct and CI is already green. I have **\*\*not\*\*** posted an LGTM on the PR yet given the dead-code finding.

Want me to push the #1 cleanup (and optionally make #2's imports lazy) to the branch, or post these as review comments on the PR?

### 34. user (2026-06-25T16:46:02.218Z)

That PR was merged by mistake, pls do a review and post the comments on the same PR and the author will work on them in the next PR

### 35. assistant (2026-06-25T16:47:13.257Z)

Understood ? I'll do a deeper adversarial review pass and post the findings as comments on

[#80](https://github.com/avidullu/muneem/pull/80). First, let me reconcile the merge state and pin a stable reference to the pre-split original for comparison.

### 36. user (2026-06-25T16:47:17.739Z)

```
=== branches containing ec091d5 ===
  origin/HEAD -> origin/main
  origin/main
  origin/refactor/split-cli
=== is ec091d5 a merge? parents ===
ec091d5886d27a0541d3d456ce091cf899ae4233 71782fd2baf4805df47e8966decff809767329f8
=== original cli.py line count at ec091d5~1 ===
611
=== new cli.py line count at ec091d5 ===
7
=== origin/main HEAD ===
ec091d5 refactor: split cli.py (~613 lines) into cli/ subpackage
71782fd Merge branch 'refactor/move-experimental-ingestion'
4bfbb68 Merge branch 'refactor/safe-amount-rename'
```

### 37. assistant (2026-06-25T16:48:07.672Z)

Confirmed: the PR was merged into `main` (now `origin/main` HEAD `ec091d5`, single-parent), so the on-disk files are the merged content and `71782fd:src/muneem/cli.py` is the stable pre-split original (611 lines).

Since ultracode is on, I'll run an adversarial multi-dimension review workflow ? independent reviewers per lens, each finding verified by a skeptic, plus a completeness critic ? so what I post is exhaustive and high-confidence.

### 38. user (2026-06-25T16:48:36.006Z)

[Request interrupted by user]

### 39. user (2026-06-25T16:48:54.893Z)

I guess since this is a refactor then just post your comments and ask them to be fixed along with the next PR

### 40. assistant (2026-06-25T16:49:23.377Z)

Got it ? ultracode off, skip the workflow. I already have the review; I'll post the findings as a consolidated follow-up comment on the merged PR. Let me write the comment body and post it.

### 41. assistant (2026-06-25T16:49:34.352Z)

## Post-merge review ? follow-ups for a later PR

The split itself is faithful: I diffed each new `cli/` module against the pre-split `cli.py` (`71782fd`) and confirmed every command body is carried over verbatim, all 12 commands are still registered (`info`, `ingest`, `eval`, `offers list`, `txns list`, `db {status,encrypt,rotate-key,backup,setup-recovery,recover}`), and the entry point `muneem.cli:app` still resolves through the thin re-export. CI was green (Ubuntu 3.11/3.12/3.13 + Windows 3.13, 478 tests, ruff + mypy). It merged before review, so logging these as cleanups to fold into the next PR ? none are blocking.

### 1. ? should-fix ? dead duplicate imports + orphaned comment in `cli/\_\_init\_\_.py`

```
```python
60 # Command registrations ? registered explicitly to avoid circular decorator imports.
61 from muneem.cli.eval import eval_statements # noqa: E402
62 from muneem.cli.ingest import ingest # noqa: E402
63
```

```

64 app.command()(ingest)
65 app.command("eval")(eval_statements)
66 # side-effects at import time. E402 is suppressed because the imports must follow
67 # the app/sub-app definitions above.
68 from muneem.cli.eval import eval_statements # noqa: E402, F401 <- dead
69 from muneem.cli.ingest import ingest      # noqa: E402, F401 <- dead
...

```

Lines 68?69 re-import `eval\_statements`/`ingest` that are already imported (61?62) and already used (64?65) ? pure no-ops, which is why they needed `# noqa: F401`. The comment on 66?67 reads mid-sentence ("side-effects at import time?"); its opening was clearly dropped in an edit. Ruff doesn't catch this (F811 doesn't fire because the first binding is *\*used\** before the redefinition, and F401 is suppressed), so it sailed through green CI. **\*\*Fix: delete lines 66?69.\*\***

## 2. ? nit ? `cli/eval.py` diverges from the stated "lazy imports" rule

The PR description lists *"Lazy imports: `extract()` imported inside functions to preserve test monkeypatching"* as a design decision, and `cli/ingest.py` honors it (imports stay inside the functions). But `cli/eval.py:9-15` hoists `extract`, `get\_parser`, `resolve\_password`, `is\_encrypted` to *\*\*module level\*\**.

It's safe today ? `tests/test\_eval\_cli.py` runs real synthetic PDFs and never patches `extract`. But it's a latent trap: an eval test that patches `muneem.extract.extract` the way `tests/test\_cli.py:39` does for ingest would silently have no effect, because `cli.eval.extract` was already bound at import time. Either make eval's imports lazy for consistency, or add a one-line comment that eval intentionally binds eagerly.

## 3. ? cosmetic ? top-level `--help` command order changed

Registration order went `info, eval, ingest` ? `info, ingest, eval` (`app.command()(ingest)` is now registered before `app.command("eval")(eval\_statements)`), so `muneem --help` lists them in a different order. No functional impact ? flag only if you want help output stable.

---

Docs: the `code-health-2026-06-25.md` tracker update is accurate and benign. Happy to send a one-line PR for #1 if that's easier than carrying it.

### 42. user (2026-06-25T16:50:57.820Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\pr80-review.md (file state is current in your context ? no need to Read it back)

### 43. user (2026-06-25T16:51:03.031Z)

<https://github.com/avidullu/muneem/pull/80#issuecomment-4802089651>

### 44. assistant (2026-06-25T16:51:06.857Z)

Posted. Comment is live: [#80 (comment)](<https://github.com/avidullu/muneem/pull/80#issuecomment-4802089651>).

### 45. user (2026-06-25T16:51:07.729Z)

cleaned up

### 46. assistant (2026-06-25T16:51:16.191Z)

Done. I posted a consolidated post-merge review on [muneem#80](<https://github.com/avidullu/muneem/pull/80#issuecomment-4802089651>) for the author to fold into the next PR.

**\*\*What it says\*\*** ? the split is faithful (every command body verbatim, all 12 commands

registered, entry point intact, CI was green), with three non-blocking follow-ups:

1. **?? should-fix??** ? dead duplicate imports + an orphaned mid-sentence comment in ``cli/___init__.py:66-69`` (ruff can't catch it; delete those 4 lines).
2. **?? nit??** ? ``cli/eval.py`` hoists ``extract`/`get_parser`/etc.` to module level, breaking the PR's own stated "lazy imports for monkeypatching" rule that ``ingest.py`` follows; safe today but a latent test trap.
3. **?? cosmetic??** ? ``--help`` command order shifted (``info, eval, ingest`` ? ``info, ingest, eval``).

I also offered to send a one-line PR for #1 if that's easier than carrying it.

Side note for completeness: the merge landed on ``main`` as ``ec091d5``, and your local ``muneem`` checkout is on ``refactor/split-cli`` pointing at that same commit ? so the branch is now fully merged and safe to delete locally/remotely whenever you like.

#### 47. user (2026-06-25T16:55:08.178Z)

can you keep polling for the PRs on the muneem repo and review it

#### 48. user (2026-06-25T16:55:11.260Z)

[Request interrupted by user]

#### 49. user (2026-06-25T16:57:28.202Z)

can you review #81

#### 50. assistant (2026-06-25T16:57:33.506Z)

Let me pull the muneem repo and gather PR #81's metadata, CI status, and diff.

#### 51. user (2026-06-25T16:57:37.998Z)

```
{"additions":5,"author":{"id":"MDQ6VXNlcjE2NzQ4Nzc=","is_bot":false,"login":"avidullu","name":"A vi"},"baseRefName":"main","body":"Addresses the post-merge review comments from PR #80.\n\n###  
1. ? Dead duplicate imports + orphaned comment in `cli/___init__.py`\nLines 66-69 were leftover from an edit ? they re-imported `eval_statements` and `ingest` that were already imported (lines 61-62) and already used (lines 64-65). The orphaned comment read mid-sentence because its opening was dropped. Deleted all four lines.\n\n### 2. ? Lazy imports in `cli/eval.py`\nMoved `extract`, `get_parser`, `resolve_password`, `is_encrypted` from module-level to lazy imports inside `_eval_statement()`, consistent with `cli/ingest.py` and preserving test monkeypatchability.\n\n### Verification\n- ? ruff check ? clean\n- ? ruff format --check ? clean\n- ? mypy ? 41 source files, no issues\n- ? pytest ? 478 passed","changedFiles":2,"deletions":8,"headRefName":"refactor/split-db-plus-review-fixes","isDraft":false,"mergeable":"MERGEABLE","number":81,"reviewDecision":"","state":"OPEN","title":"fix: address PR #80 post-merge review comments","url":"https://github.com/avidullu/muneem/pull/81"}
```

#### 52. user (2026-06-25T16:57:40.200Z)

```
=== CI checks ===  
ocr      skipping      0  
https://github.com/avidullu/muneem/actions/runs/28186627388/job/83490126067  
test (ubuntu-latest, 3.11)      pass      36s  
https://github.com/avidullu/muneem/actions/runs/28186618740/job/83490062724  
test (ubuntu-latest, 3.11)      pass      40s  
https://github.com/avidullu/muneem/actions/runs/28186627388/job/83490093260  
test (ubuntu-latest, 3.12)      pass      45s  
https://github.com/avidullu/muneem/actions/runs/28186618740/job/83490062736  
test (ubuntu-latest, 3.12)      pass      46s  
https://github.com/avidullu/muneem/actions/runs/28186627388/job/83490093091  
test (ubuntu-latest, 3.13)      pass      40s  
https://github.com/avidullu/muneem/actions/runs/28186618740/job/83490062738  
test (ubuntu-latest, 3.13)      pass      42s
```



```
https://github.com/avidullu/muneem/actions/runs/28186627388/job/83490093110
test (windows-latest, 3.13)      pass      1m25s
https://github.com/avidullu/muneem/actions/runs/28186618740/job/83490062733
test (windows-latest, 3.13)      pass      1m20s
https://github.com/avidullu/muneem/actions/runs/28186627388/job/83490093129
ocr      skipping      0
https://github.com/avidullu/muneem/actions/runs/28186618740/job/83490063391
```

```
=== files ===
src/muneem/cli/__init__.py
src/muneem/cli/eval.py
```

### 53. assistant (2026-06-25T16:57:48.301Z)

PR #81 directly addresses my two #80 comments (delete dead imports; make `eval.py` lazy). CI is green. Let me check out the branch and read the actual diff to verify the fixes are correct and complete.

### 54. user (2026-06-25T16:57:50.524Z)

```
=== full diff ===
diff --git a/src/muneem/cli/__init__.py b/src/muneem/cli/__init__.py
index 9304d52..0706b5c 100644
--- a/src/muneem/cli/__init__.py
+++ b/src/muneem/cli/__init__.py
@@ -63,7 +63,3 @@ def info() -> None:

    app.command()(ingest)
    app.command("eval")(eval_statements)
-# side-effects at import time. E402 is suppressed because the imports must follow
-# the app/sub-app definitions above.
-from muneem.cli.eval import eval_statements # noqa: E402, F401
-from muneem.cli.ingest import ingest # noqa: E402, F401
diff --git a/src/muneem/cli/eval.py b/src/muneem/cli/eval.py
index 725a4f5..12543d3 100644
--- a/src/muneem/cli/eval.py
+++ b/src/muneem/cli/eval.py
@@ -7,12 +7,8 @@
 import typer

 from muneem import db as dbmod
-from muneem.extract import extract
 from muneem.log import get_logger
-from muneem.parsers import get_parser
 from muneem.redaction import redact_filename
-from muneem.statement_unlock import resolve_password
-from muneem.unlock import is_encrypted

_log = get_logger(__name__)

@@ -23,6 +19,11 @@ def _eval_statement(pdf: Path, *, interactive: bool) -> tuple[str, int, str |
No
    Returns ``outcome, row_count, parser_name`` where ``outcome`` is the document status
    (``reconciled`` / ``unreconciled``) or a failure tag (``locked`` / ``no-parser`` /
    ``error``).
    """
+    from muneem.extract import extract
+    from muneem.parsers import get_parser
+    from muneem.statement_unlock import resolve_password
+    from muneem.unlock import is_encrypted
+
    password = ""
    if is_encrypted(pdf):
        password = resolve_password(pdf, interactive=interactive) or ""
```

## 55. assistant (2026-06-25T16:58:43.993Z)

The diff is exactly the two fixes I flagged. Let me check out the branch and read the final state of both files to confirm they're clean and the lazy-import set is precisely right.

## 56. user (2026-06-25T16:58:47.218Z)

Your branch is up to date with 'origin/refactor/split-db-plus-review-fixes'.

Already on 'refactor/split-db-plus-review-fixes'

Already up to date.

=== HEAD ===

f0b16df fix: address PR #80 post-merge review comments

## 57. user (2026-06-25T16:58:52.801Z)

```
1      """The ``muneem`` command-line interface.
2
3      Subcommands are split across focused modules: ``ingest``, ``offers``, ``txns``,
4      ``eval``, and ``db`` ? each in its own file under ``cli/``. This module assembles
5      the Typer app.
6
7      ``fetch`` (Gmail pull) lands with Phase F.
8      """
9
10     from __future__ import annotations
11
12     import typer
13
14     from muneem import __version__
15     from muneem.cli._shared import _version_callback
16     from muneem.config import load_config
17
18     app = typer.Typer(
19         add_completion=False,
20         help="muneem ? a local-first personal-finance bookkeeper.",
21         no_args_is_help=True,
22     )
23
24     # Sub-apps ? each wired in the module that defines its commands.
25     from muneem.cli.db import db_app # noqa: E402
26     from muneem.cli.offers import offers_app # noqa: E402
27     from muneem.cli.txns import txns_app # noqa: E402
28
29     app.add_typer(offers_app, name="offers")
30     app.add_typer(txns_app, name="txns")
31     app.add_typer(db_app, name="db")
32
33
34     @app.callback()
35     def main(
36         version: bool = typer.Option(
37             False,
38             "--version",
39             callback=_version_callback,
40             is_eager=True,
41             help="Show the muneem version and exit.",
42         ),
43     ) -> None:
44         """muneem command-line interface."""
45
46
47     @app.command()
48     def info() -> None:
49         """Show resolved local paths (no secrets, no sensitive data)."""
50         cfg = load_config()
```

```

51     present = "present" if cfg.config_file.is_file() else "absent"
52     typer.echo(f"muneem {__version__}")
53     typer.echo(f"config dir : {cfg.config_dir}")
54     typer.echo(f"data dir   : {cfg.data_dir}")
55     typer.echo(f"db path   : {cfg.db_path}")
56     typer.echo(f"staging   : {cfg.staging_dir}")
57     typer.echo(f"config file: {cfg.config_file} ({present})")
58
59
60     # Command registrations ? registered explicitly to avoid circular decorator imports.
61     from muneem.cli.eval import eval_statements # noqa: E402
62     from muneem.cli.ingest import ingest # noqa: E402
63
64     app.command()(ingest)
65     app.command("eval")(eval_statements)
66

```

#### 58. user (2026-06-25T16:58:53.411Z)

```

1     """The ``eval`` command ? evaluate parsers against a folder of statements."""
2
3     from __future__ import annotations
4
5     from pathlib import Path
6
7     import typer
8
9     from muneem import db as dbmod
10    from muneem.log import get_logger
11    from muneem.redaction import redact_filename
12
13    _log = get_logger(__name__)
14
15
16    def _eval_statement(pdf: Path, *, interactive: bool) -> tuple[str, int, str | None]:
17        """Run one statement through the full parse pipeline into an in-memory DB.
18
19        Returns ``(outcome, row_count, parser_name)`` where ``outcome`` is the document
status
20        (``reconciled`` / ``unreconciled``) or a failure tag (``locked`` / ``no-parser`` /
``error``).
21        """
22        from muneem.extract import extract
23        from muneem.parsers import get_parser
24        from muneem.statement_unlock import resolve_password
25        from muneem.unlock import is_encrypted
26
27        password = ""
28        if is_encrypted(pdf):
29            password = resolve_password(pdf, interactive=interactive) or ""
30            if not password:
31                return ("locked", 0, None)
32        try:
33            ex = extract(pdf, password=password)
34            parser = get_parser(ex.text)
35        except Exception as exc:

```

#### 59. assistant (2026-06-25T16:59:15.241Z)

Both files are clean. The fix is precisely scoped ? let me confirm what the green CI actually enforces (esp. a coverage gate) so my sign-off states accurately what's verified.

#### 60. user (2026-06-25T16:59:26.572Z)

```

=== CI workflow test steps ===
1:name: ci
24:      # path handling, the sqlcipher3 win_amd64 wheel). Linux covers the 3.11/3.12
spread.
34:      - name: Install
35:        run: |
38:      - name: Lint
39:        run: ruff check .
40:      - name: Format check
41:        run: ruff format --check .
42:      - name: Type check
43:        run: mypy
44:      - name: Test
45:        run: pytest --cov=muneem --cov-report=term-missing --cov-fail-under=80
59:      - name: Install Tesseract
60:        run: sudo apt-get update && sudo apt-get install -y tesseract-ocr
61:      - name: Install
62:        run: |
65:      - name: OCR tests
66:        run: pytest -m ocr

=== pyproject coverage/addopts ===
37:      "pytest-cov>=5.0",
86:addopts = "-ra -m 'not ocr and not llm and not local_statements'"

```

## 61. assistant (2026-06-25T17:00:06.122Z)

Reviewed. Here's my verdict on [#81](https://github.com/avidullu/muneem/pull/81).

## ? LGTM ? both fixes are correct, complete, and faithful

The PR does exactly what my #80 comments asked, nothing more:

**\*\*Fix 1 ? `cli/\_\_\_init\_\_.py` (dead code removal).\*\*** Lines 66?69 (the two no-op re-imports + the orphaned mid-sentence comment) are gone. The file now ends cleanly at the two registrations; the remaining late imports (61?62) are still used (64?65), so the `# noqa: E402` is still correct and there's no F401/F811. Provably a zero-behavior deletion.

**\*\*Fix 2 ? `cli/eval.py` (lazy imports).\*\*** This is the part worth getting right, and they did. The four imports moved into `\_eval\_statement` are exactly the monkeypatch-sensitive `from x import name` bindings ? `extract`, `get\_parser`, `resolve\_password`, `is\_encrypted`. Those bind the name directly, so they *must* be re-resolved at call time to see a patched `muneem.extract.extract` (the way `tests/test\_cli.py:39` patches it for ingest). `dbmod` was correctly **\*\*left\*\*** at module level: it's a `import ? as` module alias, so `dbmod.connect` resolves dynamically and stays monkeypatch-safe regardless. This restores the original pre-split `\_eval\_statement` behavior and matches `ingest.py`. Precisely scoped ? no over-correction.

**\*\*No new issues, no production behavior change.\*\*** The only observable delta is restored test-monkeypatchability, which is the intent. My #3 (cosmetic `--help` order) was explicitly optional and correctly left alone.

**\*\*What green CI verifies here\*\*** (I relied on it rather than re-running): `ruff check`, `ruff format --check`, `mypy`, 478 tests, **\*\*and `--cov-fail-under=80`\*\*** ? all passing on Ubuntu 3.11/3.12/3.13 + Windows 3.13. That covers the full pre-merge gate; coverage can't have regressed since the change only relocates always-executed import lines.

**\*\*One non-blocking heads-up:\*\*** the branch is named `refactor/split-db-plus-review-fixes`, but this PR contains *only* the two review fixes ? no `db.py` split. The title/body are accurate; just flagging in case a db split was meant to ride along and got dropped.

---

Want me to **\*\*post this LGTM on #81\*\***? I held off auto-posting since you said "review" (not "post") this time, and an LGTM may prompt you to merge. And do you still want me to **\*\*keep**

polling muneem for new PRs\*\* and review them as they land ? if so, at what cadence?

**62. user (2026-06-25T17:00:33.039Z)**

yes pls post an LGTM to it

**63. assistant (2026-06-25T17:00:43.584Z)**

## LGTM ?

Both fixes from the #80 follow-up are correct, complete, and faithful ? nothing over-corrected.

**\*\*1. `cli/\_\_\_init\_\_.py` ? dead code removal.\*\*** Lines 66?69 (the two no-op re-imports + the orphaned mid-sentence comment) are gone. The remaining late imports (61?62) are still used at 64?65, so the `# noqa: E402` stays valid and there's no F401/F811. Provably a zero-behavior deletion.

**\*\*2. `cli/eval.py` ? lazy imports.\*\*** Precisely scoped: the four imports moved into `\_eval\_statement` are exactly the monkeypatch-sensitive `from X import name` bindings (`extract`, `get\_parser`, `resolve\_password`, `is\_encrypted`) ? those bind the name directly and must be re-resolved at call time to see a patched `muneem.extract.extract` (the way `tests/test\_cli.py:39` patches it for ingest). `dbmod` was correctly **\*\*left\*\*** at module level: it's an `import` ? as `alias`, so `dbmod.connect` resolves dynamically and stays monkeypatch-safe. This restores the original pre-split `\_eval\_statement` behavior and matches `ingest.py`.

No new issues, no production behavior change ? the only observable delta is restored test-monkeypatchability, which is the intent.

**\*\*Verified via green CI\*\*** (relied on, not re-run): `ruff check`, `ruff format --check`, `mypy`, 478 tests, **\*\*and `--cov-fail-under=80`\*\*** all passing on Ubuntu 3.11/3.12/3.13 + Windows 3.13. That's the full pre-merge gate; coverage can't regress since the change only relocates always-executed import lines.

**\*\*Non-blocking heads-up:\*\*** the branch is named `refactor/split-db-plus-review-fixes`, but this PR contains only the two review fixes ? no `db.py` split. Title/body are accurate; just flagging in case a db split was meant to ride along and got dropped.

**64. user (2026-06-25T17:00:43.618Z)**

File created successfully at: C:\Users\avidu\AppData\Local\Temp\pr81-lgtm.md (file state is current in your context ? no need to Read it back)

**65. user (2026-06-25T17:00:48.998Z)**

<https://github.com/avidullu/muneem/pull/81#issuecomment-4802168838>

**66. assistant (2026-06-25T17:00:55.541Z)**

Posted: [#81 LGTM comment](<https://github.com/avidullu/muneem/pull/81#issuecomment-4802168838>).

Written LGTM (not a formal GitHub Approve, per our convention for your own account). The fix cleanly closes out both #80 findings, CI's green across the full gate (lint/format/mypy/478 tests/coverage ?80%), and it's safe to merge whenever you like.

Still open from my last message: do you want me to **\*\*keep polling muneem for new PRs\*\*** and review them as they land? If yes, tell me a cadence (e.g. every 15?30 min) and whether to **\*\*auto-post\*\*** reviews or just surface them here for you to approve first.

**67. user (2026-06-25T17:04:00.909Z)**

#82 is up for review. Can you pls do it

**68. user (2026-06-25T17:04:07.875Z)**

```
{
  "additions": 581,
  "author": {
    "id": "MDQ6VXNlcjE2NzQ4Nzc=",
    "is_bot": false,
    "login": "avidullu",
    "name": "Avi"
  },
  "baseRefName": "main",
  "body": "Splits the monolithic `db.py` (~590 lines) into a `db/` subpackage with focused modules.\n\n### Before\n```\nsrc/muneem/db.py (590 lines ? schema DDL, key helpers, connections, crypto, queries all in one file)\n```\n\nAfter\n```\nsrc/muneem/db/\n  __init__.py (35 lines ? re-exports all public names, backward-compatible)\n  schema.py (105 lines ? _SCHEMA DDL, SCHEMA_VERSION, init_db)\n  connection.py (155 lines ? key validation, is_plaintext_sqlite, connect, connect_encrypted)\n  crypto.py (100 lines ? encrypt_database, rekey_database, backup_database)\n  queries.py (180 lines ? document CRUD, provenance, offers, per-bank inserts + listings)\n```\n\n### Key design decisions\n- **`connection.py`** not `connect.py` ? avoids clashing with the `connect()` function when using `import muneem.db.connect` (Python resolves the package attribute first)\n- **All public names re-exported** from `__init__.py` ? every `from muneem import db as dbmod` import continues to work\n- **Private symbols stay in submodules** ? `_SCHEMA`, `_open_encrypted` are not re-exported; tests import from the specific submodule or monkeypatch at the module level\n\n### Verification\n- ? `ruff check` ? clean\n- ? `ruff format --check` ? clean\n- ? `mypy` ? 45 source files, no issues\n- ? `pytest` ? 478 passed\n- ? `muneem --version` ? works\n\nFixes recommendation #6 from the code-health report (PR #74).",
  "changedFiles": 8,
  "deletions": 585,
  "headRefName": "refactor/split-db",
  "isDraft": false,
  "mergeable": "MERGEABLE",
  "number": 82,
  "reviewDecision": "",
  "state": "OPEN",
  "title": "refactor: split db.py (~590 lines) into db/ subpackage",
  "url": "https://github.com/avidullu/muneem/pull/82"
}
```

## 69. user (2026-06-25T17:04:10.365Z)

```
=== CI checks ===
test (windows-latest, 3.13)    pending 0
https://github.com/avidullu/muneem/actions/runs/28187001914/job/83491399767
test (ubuntu-latest, 3.11)     pass    34s
https://github.com/avidullu/muneem/actions/runs/28187001914/job/83491399778
test (ubuntu-latest, 3.12)     pass    49s
https://github.com/avidullu/muneem/actions/runs/28187001914/job/83491399735
test (ubuntu-latest, 3.13)     pass    43s
https://github.com/avidullu/muneem/actions/runs/28187001914/job/83491399727
ocr    skipping                0
https://github.com/avidullu/muneem/actions/runs/28187001914/job/83491401220
test (ubuntu-latest, 3.11)     pass    39s
https://github.com/avidullu/muneem/actions/runs/28186994783/job/83491375380
test (ubuntu-latest, 3.12)     pass    51s
https://github.com/avidullu/muneem/actions/runs/28186994783/job/83491375393
test (ubuntu-latest, 3.13)     pass    36s
https://github.com/avidullu/muneem/actions/runs/28186994783/job/83491375328
test (windows-latest, 3.13)    pass    1m30s
https://github.com/avidullu/muneem/actions/runs/28186994783/job/83491375348
ocr    skipping                0
https://github.com/avidullu/muneem/actions/runs/28186994783/job/83491376320
```

```
=== changed files ===
src/muneem/db.py
src/muneem/db/__init__.py
src/muneem/db/connection.py
src/muneem/db/crypto.py
src/muneem/db/queries.py
src/muneem/db/schema.py
tests/test_db_crypto_ops.py
tests/test_migrations.py
```

## 70. assistant (2026-06-25T17:04:38.240Z)

PR #82 is the `db.py` split that the #81 branch name hinted at ? and this one touches `**crypto/connection/key-handling code**`, so I'll verify the moved logic is byte-faithful, not just that it imports. Note: `**Windows CI is still pending**` (Ubuntu 3.11/3.12/3.13 green) ? Windows is the platform that exercises the `sqlcipher3` wheel, so I'll factor that in. Let me pull the original and the new modules without disturbing the working tree.

## 71. user (2026-06-25T17:04:40.946Z)

```

merge-base: ec12b59eb1320061c7fbfdefdcc8133a95970a61
=== original db.py lines ===
583
=== public symbols in original db.py (def/class/CONST) ===
21:_log = get_logger(__name__)
23:SCHEMA_VERSION = 4
27:_SQLITE_MAGIC = b"SQLite format 3\x00"
29:_SCHEMA = ""
113:def _validate_raw_key(key: str) -> str:
128:def _raw_key_pragma(key: str) -> str:
133:def is_plaintext_sqlite(db_path: Path | str) -> bool:
146:def _harden_permissions(path: str) -> None:
156:def _open_encrypted(path: str, key: str) -> sqlite3.Connection:
184:def connect(db_path: Path | str, *, key: str | None = None) -> Iterator[sqlite3.Connection]:
209:def key_opens_database(db_path: Path | str, key: str) -> bool:
219:def connect_encrypted(db_path: Path | str) -> Iterator[sqlite3.Connection]:
272:def encrypt_database(db_path: Path | str, key: str) -> int:
335:def rekey_database(db_path: Path | str, current_key: str, new_key: str) -> None:
352:def backup_database(db_path: Path | str, dest: Path | str, key: str) -> None:
396:def init_db(conn: sqlite3.Connection) -> None:
407:def get_document_by_sha(conn: sqlite3.Connection, sha256: str) -> sqlite3.Row | None:
411:def clear_document_rows(conn: sqlite3.Connection, document_id: int) -> None:
428:def _insert_row(conn: sqlite3.Connection, table: str, fields: dict) -> int:
442:def insert_document(conn: sqlite3.Connection, **fields: object) -> int:
447:def record_provenance(
475:def record_parse(
498:def get_provenance(conn: sqlite3.Connection, document_id: int) -> sqlite3.Row | None:
502:def insert_offer(conn: sqlite3.Connection, **fields: object) -> int:
506:def list_offers(conn: sqlite3.Connection) -> list[sqlite3.Row]:
518:def insert_axis_transaction(conn: sqlite3.Connection, **fields: object) -> int:
522:def insert_axis_cc_transaction(conn: sqlite3.Connection, **fields: object) -> int:
528:def list_axis_cc_transactions(conn: sqlite3.Connection, limit: int = 50) ->
list[sqlite3.Row]:
536:def list_axis_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
544:def insert_hdfc_transaction(conn: sqlite3.Connection, **fields: object) -> int:
548:def list_hdfc_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
556:def insert_au_transaction(conn: sqlite3.Connection, **fields: object) -> int:
560:def list_all_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
578:def list_au_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:

```

## 72. user (2026-06-25T17:04:42.651Z)

```

=== test file diffs ===
diff --git a/tests/test_db_crypto_ops.py b/tests/test_db_crypto_ops.py
index 9c3d367..378b372 100644
--- a/tests/test_db_crypto_ops.py
+++ b/tests/test_db_crypto_ops.py
@@ -12,6 +12,7 @@
 from typer.testing import CliRunner # noqa: E402

+import muneem.db.connection as dbconnect_mod # noqa: E402
 from muneem import db as dbmod # noqa: E402
 from muneem import db_key # noqa: E402
 from muneem.cli import app # noqa: E402
@@ -161,7 +162,7 @@ def test_encrypt_database_cleans_temp_on_failure(tmp_path, monkeypatch):
     def boom(*_a, **_k): # fail the post-export verification open
         raise RuntimeError("boom during verify")

-    monkeypatch.setattr(dbmod, "_open_encrypted", boom)
+    monkeypatch.setattr(dbconnect_mod, "_open_encrypted", boom)
     with pytest.raises(RuntimeError):
         dbmod.encrypt_database(p, KEY_A)
     assert not (tmp_path / "led.db.enc.tmp").exists() # partial temp removed
diff --git a/tests/test_migrations.py b/tests/test_migrations.py

```

```

index de0d4f2..daf4bfb 100644
--- a/tests/test_migrations.py
+++ b/tests/test_migrations.py
@@ -6,9 +6,9 @@

from muneem import db as dbmod
from muneem import migrations
+from muneem.db.schema import _SCHEMA as _BASE
from muneem.errors import MigrationError

-_BASE = dbmod._SCHEMA
+_VER = dbmod.SCHEMA_VERSION # 4

```

### 73. assistant (2026-06-25T17:05:25.831Z)

The test changes look exactly right at a glance (patch `\_open\_encrypted` where it now lives in `connection.py`; import private `\_SCHEMA` from the submodule). The critical thing for a crypto refactor is that **no function body changed** during the move. Let me prove that with an AST-level comparison of every symbol ? original `db.py` vs the merged `db/` modules.

### 74. user (2026-06-25T17:05:29.293Z)

```

files staged in /tmp/tmp.A1gmiP4Pyp:
total 2104
drwxr-xr-x 1 avidu 197609    0 Jun 25 22:35 .
drwxr-xr-x 1 avidu 197609    0 Jun 25 22:35 ..
-rw-r--r-- 1 avidu 197609 1357 Jun 25 22:35 __init__.py
-rw-r--r-- 1 avidu 197609 5080 Jun 25 22:35 connection.py
-rw-r--r-- 1 avidu 197609 3995 Jun 25 22:35 crypto.py
-rw-r--r-- 1 avidu 197609 23071 Jun 25 22:35 orig.py
-rw-r--r-- 1 avidu 197609 5990 Jun 25 22:35 queries.py
-rw-r--r-- 1 avidu 197609 2824 Jun 25 22:35 schema.py

```

### 75. user (2026-06-25T17:05:53.498Z)

```

Exit code 1
Traceback (most recent call last):
  File "C:\Users\avidu\AppData\Local\Temp\astcmp.py", line 2, in <module>
    TMP = open('/tmp/cmpdir.txt').read().strip()
    ~~~~~^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
FileNotFoundError: [Errno 2] No such file or directory: '/tmp/cmpdir.txt'

```

### 76. assistant (2026-06-25T17:06:14.146Z)

The Windows Python can't resolve git-bash `/tmp` paths. Let me re-stage and run the comparison passing Windows paths.

### 77. user (2026-06-25T17:06:23.071Z)

```

orig symbols=33 new symbols=34 AST-identical=11
MISSING: NONE
CHANGED(AST differs): ['_SCHEMA', '_harden_permissions', '_open_encrypted', '_raw_key_pragma',
'_validate_raw_key', 'backup_database', 'clear_document_rows', 'connect', 'connect_encrypted',
'encrypt_database', 'init_db', 'insert_document', 'is_plaintext_sqlite', 'key_opens_database',
'list_all_transactions', 'list_auth_transactions', 'list_axis_cc_transactions',
'list_axis_transactions', 'list_hdfc_transactions', 'record_parse', 'record_provenance',
'rekey_database']
ADDED: ['_JOIN_STATUS']

```

### 78. assistant (2026-06-25T17:06:59.442Z)

22 symbols differ at the AST level and a new `\_JOIN\_STATUS` constant appeared ? so this is **not** a pure mechanical move; some bodies were also edited (likely cross-module call re-qualification + a SQL dedup). That's exactly what needs scrutiny in a "refactor" PR. Let me get



the actual per-function diffs.

## 79. user (2026-06-25T17:07:08.394Z)

```
===== _validate_raw_key =====
-   """Require exactly 64 hex chars before splicing key material into a PRAGMA/ATTACH string.
-
-   Anything else would either silently switch SQLCipher from raw-key mode to passphrase-KDF
-   mode (a *different* encryption that still "works" ? and then never opens again with the
-   real raw key), or break out of the SQL literal. Refuse loudly instead.
-   """
+   """Require exactly 64 hex chars before splicing key material into a PRAGMA/ATTACH
string."""

===== _raw_key_pragma =====
-   """The SQLCipher raw-key PRAGMA: use the 256-bit ``key`` (64 hex chars) directly, no
KDF."""
+   """The SQLCipher raw-key PRAGMA: use the 256-bit key (64 hex chars) directly, no KDF."""

===== is_plaintext_sqlite =====
-   """True iff the file exists and begins with the *unencrypted* SQLite magic header.
-
-   Lets us refuse to silently treat a plaintext ledger as encrypted (and vice-versa)
-   without needing the key.
-   """
+   """True iff the file exists and begins with the unencrypted SQLite magic header."""

===== _harden_permissions =====
-   """Best-effort: restrict the DB file to its owner (0600). No-op on failure or
``:memory:``."""
+   """Best-effort: restrict the DB file to its owner (0600). No-op on failure or :memory:."""

===== _open_encrypted =====
-   """Open an encrypted connection via SQLCipher, verifying the key before returning.
-
-   A keyed open NEVER silently downgrades to plaintext: if the binding is missing, or the
-   key is wrong (or the file is not an encrypted muneem ledger), this raises.
-   """
+   """Open an encrypted connection via SQLCipher, verifying the key before returning."""
-   except ImportError as exc: # pragma: no cover - exercised only without the binding
+   except ImportError as exc: # pragma: no cover
-   conn.execute(_raw_key_pragma(key)) # MUST be the first statement on the connection
+   conn.execute(_raw_key_pragma(key))
-   conn.execute("SELECT count(*) FROM sqlite_master") # forces key verification
+   conn.execute("SELECT count(*) FROM sqlite_master")

===== connect =====
-   With ``key`` (a 64-hex-char SQLCipher key) the database is **encrypted at rest** via
-   SQLCipher. Without it, a plaintext stdlib :mod:`sqlite3` connection is used ? the path for
-   ``:memory:`` and tests. A keyed open never silently falls back to plaintext (see
-   :func:`_open_encrypted`); the production ledger is opened via :func:`connect_encrypted`.
+   With ``key`` the database is encrypted at rest via SQLCipher. Without it, plaintext
+   stdlib sqlite3 is used ? the path for :memory: and tests.

===== key_opens_database =====
-   """True iff ``key`` opens the encrypted ledger at ``db_path`` (verification only, no
reads)."""
+   """True iff ``key`` opens the encrypted ledger at ``db_path`` (verification only)."""

===== connect_encrypted =====
-   The production path: resolves (or mints, on first use) the 256-bit key from the OS
-   keychain and opens with SQLCipher. If ``db_path`` is a pre-existing *plaintext* SQLite
-   file, refuses with a clear pointer to ``muneem db encrypt`` rather than a cryptic
-   "file is not a database" error.
-
```

```

- Self-healing for interrupted key rotations: rotation stages the incoming key under a
- *pending* keychain entry before the file is rekeyed (see ``cli rotate-key``). If the
- primary key no longer opens the file but the pending key does, the rotation completed
- on the file side and was interrupted before promotion ? promote the pending key and
- carry on. With no pending key (or one that doesn't fit), the original wrong-key error
- propagates untouched.
+ Self-healing for interrupted key rotations: if the primary key no longer opens the file
+ but the pending key does, promote it automatically.
- # Mint a key only for a brand-new ledger. If an encrypted ledger already exists, a missing
- # keychain entry means the key was lost (keychain wiped, new machine) ? minting a fresh key
- # could never open it and would only bury the problem; point at recovery instead.
- # Raises the same clear wrong-key error if the pending key doesn't fit either.
- db_key.set_key(pending) # heal: the file was already rekeyed ? promote
+ db_key.set_key(pending)

```

==== encrypt\_database =====

```

- """Migrate a plaintext SQLite ledger to an encrypted one in place, via
- ``sqlcipher_export``.
+ """Migrate a plaintext SQLite ledger to an encrypted one in place, via sqlcipher_export.
- The original is preserved as ``<name>.plaintext.bak`` for the caller to verify and then
- securely delete (it still holds the cleartext data). Returns the number of user tables
- copied. Raises :class:`EncryptedDatabaseError` if ``db_path`` is not a plaintext SQLite
file.
+ The original is preserved as ``<name>.plaintext.bak``. Returns the number of user tables
+ copied.
- src = sqlcipher3.connect(str(path)) # open the plaintext source (no key)
+ src = sqlcipher3.connect(str(path))
- # SQL string literal: forward slashes; single quotes doubled (a path like
- # "ledger.db" must not break out of the literal).
- # Verify the encrypted copy BEFORE the destructive swap: it must open with the key,
pass
- # an integrity check, and carry every table across. (The plaintext .bak is kept either
way,
- # so a failed check is recoverable ? but we refuse to swap in a bad file.)
- enc_tmp.rename(path) # consumed on success
+ enc_tmp.rename(path)
- # Clean the temp ONLY while the original ledger is still at `path` (pre-swap failure,
- # or success ? where the rename already consumed the temp). If the swap moved the
- # original aside but the final rename failed, enc_tmp is the only VERIFIED encrypted
- # copy ? keep it.

```

==== rekey\_database =====

```

- """Re-encrypt the ledger under ``new_key`` in place (SQLCipher ``PRAGMA rekey``).
-
- Opens with ``current_key`` first (which verifies it), then rekeys. Raises
- :class:`EncryptedDatabaseError` if ``current_key`` is wrong or the file is not an encrypted
- muneem ledger. The caller is responsible for updating the keychain entry to ``new_key``
only
- after this returns successfully.
- """
- conn = _open_encrypted(str(db_path), current_key) # validates current_key
+ """Re-encrypt the ledger under ``new_key`` in place (SQLCipher PRAGMA rekey)."""
+ conn = _open_encrypted(str(db_path), current_key)

```

==== backup\_database =====

```

- """Write an **encrypted** online-backup copy of the ledger to ``dest`` (same key).
-
- Uses the SQLite Online Backup API for a consistent snapshot ? safe even with WAL active, so
no
- naive file copy and no manual checkpoint needed. The destination is keyed with the same
key, so
- it stays encrypted at rest. Raises :class:`EncryptedDatabaseError` if ``key`` is wrong.
-
- Destructive-step ordering is deliberate: the destination is refused if it is the live
ledger

```

```

-   itself, the key is validated against the source *before* the destination is touched, and
the
-   snapshot is written to a temp file that atomically replaces ``dest`` only on success ? so
-   neither the ledger nor a previous good backup can be destroyed by a bad invocation.
-   """"
+   """"Write an encrypted online-backup copy of the ledger to ``dest`` (same key).""""

```

#### 80. assistant (2026-06-25T17:07:31.575Z)

These are all **\*\*docstring/comment trims\*\*** ? ``ast.dump`` includes docstrings, so editorial rewrites flag as "differs" even though the executable logic is untouched. So far no logic change in the crypto/connection functions. But I still need to see the **\*\*query functions\*\*** and the new ``_JOIN_STATUS`` constant ? a SQL dedup *can* change results. Let me look at those specifically.

#### 81. user (2026-06-25T17:07:34.943Z)

```

=== NEW _JOIN_STATUS constant in queries.py ===
135:_JOIN_STATUS = "JOIN documents d ON d.id = t.document_id ORDER BY t.id DESC LIMIT ?"
136-
137-
138-def list_axis_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
139-    return conn.execute(
140-        f"SELECT t.*, d.status AS doc_status FROM axis_transactions t {_JOIN_STATUS}",
141-        (limit,),
142-    ).fetchall()
143-
144-
145-def list_axis_cc_transactions(conn: sqlite3.Connection, limit: int = 50) ->
list[sqlite3.Row]:
146-    return conn.execute(
147-        f"SELECT t.*, d.status AS doc_status FROM axis_cc_transactions t {_JOIN_STATUS}",
148-        (limit,),
149-    ).fetchall()
150-
151-
152-def list_hdfc_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
153-    return conn.execute(
154-        f"SELECT t.*, d.status AS doc_status FROM hdfc_transactions t {_JOIN_STATUS}",
155-        (limit,),
156-    ).fetchall()
157-
158-
159-def list_au_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
160-    return conn.execute(
161-        f"SELECT t.*, d.status AS doc_status FROM au_transactions t {_JOIN_STATUS}",
162-        (limit,),
163-    ).fetchall()
164-
165-
166-def list_all_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
167-    """"Cross-bank rows from the derived serving view, with each parse's confidence joined
in.""""

=== ORIGINAL list_all_transactions + list_au_transactions + a per-bank listing (for SQL
comparison) ===
def get_provenance(conn: sqlite3.Connection, document_id: int) -> sqlite3.Row | None:
    return conn.execute("SELECT * FROM provenance WHERE document_id = ?",
(document_id,)).fetchone()

def insert_offer(conn: sqlite3.Connection, **fields: object) -> int:
    return _insert_row(conn, "offers", fields)

def list_offers(conn: sqlite3.Connection) -> list[sqlite3.Row]:

```

```

return conn.execute(
    """
    SELECT o.id, o.merchant, o.promo_code, o.expiry_raw, o.description, o.source,
           d.external_id, d.path
    FROM offers o
    JOIN documents d ON d.id = o.document_id
    ORDER BY o.id
    """
).fetchall()

def insert_axis_transaction(conn: sqlite3.Connection, **fields: object) -> int:
    return _insert_row(conn, "axis_transactions", fields)

def insert_axis_cc_transaction(conn: sqlite3.Connection, **fields: object) -> int:
    return _insert_row(conn, "axis_cc_transactions", fields)

```

**Per-bank listings carry the parse verdict (documents.status as doc\_status) so a consumer can tell trusted rows from quarantined/legacy residue without a second query.**

```

def list_axis_cc_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
    return conn.execute(
        "SELECT t.*, d.status AS doc_status FROM axis_cc_transactions t "
        "JOIN documents d ON d.id = t.document_id ORDER BY t.id DESC LIMIT ?",
        (limit,),
    ).fetchall()

def list_axis_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
    return conn.execute(
        "SELECT t.*, d.status AS doc_status FROM axis_transactions t "
        "JOIN documents d ON d.id = t.document_id ORDER BY t.id DESC LIMIT ?",
        (limit,),
    ).fetchall()

def insert_hdfc_transaction(conn: sqlite3.Connection, **fields: object) -> int:
    return _insert_row(conn, "hdfc_transactions", fields)

def list_hdfc_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
    return conn.execute(
        "SELECT t.*, d.status AS doc_status FROM hdfc_transactions t "
        "JOIN documents d ON d.id = t.document_id ORDER BY t.id DESC LIMIT ?",
        (limit,),
    ).fetchall()

def insert_au_transaction(conn: sqlite3.Connection, **fields: object) -> int:
    return _insert_row(conn, "au_transactions", fields)

def list_all_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
    """Cross-bank rows from the derived serving view, with each parse's confidence joined in.

    Read-only over ``v_transactions`` (migration v6) ? the flat per-issuer tables remain the
    source of truth (DESIGN §8.7).
    """
    return conn.execute(
        """
        SELECT v.*, p.confidence
        FROM v_transactions v

```

```

        LEFT JOIN provenance p ON p.document_id = v.document_id
        ORDER BY v.document_id DESC
        LIMIT ?
        """
        (limit,),
    ).fetchall()

```

```

def list_aud_transactions(conn: sqlite3.Connection, limit: int = 50) -> list[sqlite3.Row]:
    return conn.execute(
        "SELECT t.*, d.status AS doc_status FROM aud_transactions t "
        "JOIN documents d ON d.id = t.document_id ORDER BY t.id DESC LIMIT ?",
        (limit,),
    ).fetchall()

```

## 82. assistant (2026-06-25T17:07:57.488Z)

The ``_JOIN_STATUS`` dedup is byte-identical SQL ? original ``"...t " + "JOIN documents d ON d.id = t.document_id ORDER BY t.id DESC LIMIT ?"``` reconstructs exactly via the f-string, same param binding. Now let me run the `**definitive**` check: strip docstrings and re-compare ASTs, so only genuine `*executable-logic*` changes remain.

## 83. user (2026-06-25T17:08:06.029Z)

```

EXECUTABLE-LOGIC DIFFERS (docstrings stripped): ['_SCHEMA', 'list_aud_transactions',
'list_axis_cc_transactions', 'list_axis_transactions', 'list_hdfc_transactions']
missing: NONE added: ['_JOIN_STATUS']

```

## 84. assistant (2026-06-25T17:08:17.199Z)

Excellent ? that's the definitive result. Every crypto/connection/query function is `**executable-identical**` (only docstrings/comments trimmed). Only two real deltas remain: the 4 ``list_`` functions (verified byte-identical SQL via ``_JOIN_STATUS``) and `**`_SCHEMA`**` ? the table DDL. A schema-string change is the one thing that could actually matter, so let me diff it precisely.

## 85. user (2026-06-25T17:08:22.594Z)

```

identical? False
--- orig
+++ new
@@ -30,8 +30,6 @@

-- We use separate tables for each bank (e.g. axis_transactions, hdfc_transactions)
-- instead of a unified table to preserve raw schema fidelity of the source PDFs.
--- Normalization into a unified view layer is planned for a later phase
--- once all bank schemas are stabilized.
CREATE TABLE IF NOT EXISTS axis_transactions (
    id INTEGER PRIMARY KEY,
    document_id INTEGER NOT NULL REFERENCES documents(id) ON DELETE CASCADE,

```

## 86. assistant (2026-06-25T17:08:46.797Z)

Only change to ``_SCHEMA`` is a 2-line SQL comment removal ? SQLite ignores ``--`` comments, so the DDL is semantically identical.

`**Net: every executable path is faithful.**` Now two public-surface checks remain: confirm ``__init__.py`` re-exports the entire public API (especially ``EncryptedDatabaseError``, which callers use as ``dbmod.EncryptedDatabaseError``), and re-poll the Windows CI job.

## 87. user (2026-06-25T17:08:50.105Z)

```

=== new db/__init__.py ===
"""Local SQLite store ? encrypted ledger, schema, and per-table helpers.

```

The ``db/`` package splits what was a ~590-line ``db.py`` monolith into focused modules:

- ``schema`` ? DDL + ``init\_db``
- ``connect`` ? key helpers, plaintext/encrypted connections
- ``crypto`` ? encrypt, rekey, backup
- ``queries`` ? document CRUD, provenance, offers, per-bank transactions

All public names are re-exported here, so ``from muneem.db import connect\_encrypted``  
(and all other existing imports) still work unchanged.

"""

```
from muneem.db.connection import ( # noqa: F401
    connect,
    connect_encrypted,
    is_plaintext_sqlite,
    key_opens_database,
)
from muneem.db.crypto import backup_database, encrypt_database, rekey_database # noqa: F401
from muneem.db.queries import ( # noqa: F401
    clear_document_rows,
    get_document_by_sha,
    get_provenance,
    insert_au_transaction,
    insert_axis_cc_transaction,
    insert_axis_transaction,
    insert_document,
    insert_hdfc_transaction,
    insert_offer,
    list_all_transactions,
    list_au_transactions,
    list_axis_cc_transactions,
    list_axis_transactions,
    list_hdfc_transactions,
    list_offers,
    record_parse,
    record_provenance,
)
from muneem.db.schema import SCHEMA_VERSION, init_db # noqa: F401
from muneem.errors import EncryptedDatabaseError # noqa: F401

=== where is EncryptedDatabaseError defined/imported in the split? ===
-- __init__.py --
41:from muneem.errors import EncryptedDatabaseError # noqa: F401
-- schema.py --
-- connection.py --
13:from muneem.errors import EncryptedDatabaseError
22:     raise EncryptedDatabaseError(
58:     raise EncryptedDatabaseError(
68:     raise EncryptedDatabaseError(
105: except EncryptedDatabaseError:
122:     raise EncryptedDatabaseError(
129:     raise EncryptedDatabaseError(
139: except EncryptedDatabaseError:
-- crypto.py --
16:from muneem.errors import EncryptedDatabaseError
29:     raise EncryptedDatabaseError(
56:     raise EncryptedDatabaseError(
97:     raise EncryptedDatabaseError(
-- queries.py --

=== callers using dbmod.EncryptedDatabaseError ===
src/muneem/cli/db.py:142:     except dbmod.EncryptedDatabaseError as exc:
src/muneem/cli/db.py:202:     except dbmod.EncryptedDatabaseError as exc:
src/muneem/cli/_shared.py:20:     except dbmod.EncryptedDatabaseError as exc:
Binary file src/muneem/cli/__pycache__/db.cpython-313.pyc matches
```

Binary file src/muneem/cli/\_\_pycache\_\_/\_shared.cpython-313.pyc matches

#### 88. user (2026-06-25T17:08:52.183Z)

```
=== Windows CI status now ===
test (windows-latest, 3.13)      pass      2m10s
https://github.com/avidullu/muneem/actions/runs/28187001914/job/83491399767
test (windows-latest, 3.13)      pass      1m30s
https://github.com/avidullu/muneem/actions/runs/28186994783/job/83491375348
```

#### 89. assistant (2026-06-25T17:09:26.082Z)

All green now ? Windows CI passed (2m10s), so the full gate (Ubuntu 3.11/3.12/3.13 + Windows 3.13, sqlcipher crypto tests included) is green. The `\_\_init\_\_.py` re-exports the complete public API including `EncryptedDatabaseError` (which `cli/db.py` and `cli/\_shared.py` rely on as `dbmod.EncryptedDatabaseError`). One last belt-and-suspenders check: confirm no *\*non-test\** code reaches for a now-private db internal via the package.

#### 90. user (2026-06-25T17:09:29.917Z)

```
=== any non-test references to private db internals
(_SCHEMA/_open_encrypted/_insert_row/_validate_raw_key) via db/dbmod? ===
src/muneem/migrations.py:9:- **Baseline.** The current schema (`db._SCHEMA`, version
`db.SCHEMA_VERSION`) is the *frozen*

=== does migrations.py touch db private DDL? ===
9:- **Baseline.** The current schema (`db._SCHEMA`, version `db.SCHEMA_VERSION`) is the *frozen*
31:_SCHEMA_MIGRATIONS = ""
196:     conn.executescript(_SCHEMA_MIGRATIONS) # bookkeeping table first (idempotent)
```

#### 91. assistant (2026-06-25T17:10:33.878Z)

Review complete ? and this one I verified deeply, since it moves *\*\*crypto/key-handling code\*\**. Here's the verdict on [#82](https://github.com/avidullu/muneem/pull/82).

### ? Faithful split ? behavior provably preserved, full CI now green

I didn't take "it's a refactor" on faith. I diffed every symbol at the *\*\*AST level\*\** (original `db.py` vs the merged `db/` modules), and after stripping docstrings, *\*\*all 33 symbols are executable-identical\*\** except two, both of which I verified by hand:

- *\*\*The 4 `list\_` per-bank functions\*\** were refactored to share a new `\_JOIN\_STATUS` constant. The f-string reconstructs *\*\*byte-identical SQL\*\** ? same `JOIN documents d ON d.id = t.document\_id ORDER BY t.id DESC LIMIT ?`, same `(limit,)` binding. No result change.
- *\*\*`\_SCHEMA`\*\** differs by exactly one thing: a 2-line `--` SQL comment was removed. SQLite ignores `--` comments, so *\*\*the DDL is semantically identical\*\** ? same tables, columns, constraints.

Everything else ? `encrypt\_database`, `rekey\_database`, `backup\_database`, `connect`, `connect\_encrypted`, `\_open\_encrypted`, key validation, all CRUD ? is *\*\*byte-faithful\*\** (only docstrings/comments changed).

*\*\*Public surface intact:\*\** `db/\_\_init\_\_.py` re-exports all 26 public names *\*\*plus* `EncryptedDatabaseError`, so `from muneem import db as dbmod; dbmod.X` keeps working ? including the `dbmod.EncryptedDatabaseError` that `cli/db.py` and `cli/\_shared.py` catch. Private symbols (`\_SCHEMA`, `\_open\_encrypted`, ?) are correctly *\*not\** re-exported; the *\*\*two* test edits are exactly the right, minimal follow-through ? patch `\_open\_encrypted` at its new home `muneem.db.connection`, import `\_SCHEMA` from `muneem.db.schema`. No test was weakened. No non-test code references a now-private internal.

*\*\*CI is fully green now\*\** ? Windows 3.13 finished passing (2m10s) while I was reviewing; that's the job that exercises the sqlcipher3 wheel + crypto tests. Full gate covered (ruff/format/mypy/478 tests/cov ?80%) on Ubuntu 3.11/3.12/3.13 + Windows.

## Non-blocking notes (for a follow-up, not a blocker)

1. **The PR says "pure split" but also bundles editorial trims.** Several **security-rationale comments** were dropped from the crypto/connection code ? e.g. ``_raw_key_pragma`'s` **"MUST be the first statement on the connection"**, ``_open_encrypted`'s` **"forces key verification"**, ``encrypt_database`'s` **destructive-step-ordering explanation**, and ``connect_encrypted`'s` **rekey self-healing rationale**. Behavior-neutral, but those comments documented **why the code is safe**. For security-critical code I'd lean toward keeping them ? and keeping a "split" purely mechanical also makes it trivially reviewable.

2. **Doc nit:** the ``db/___init___.py`` docstring lists the module as ``connect``, but it's actually ``connection.py`` ? ironic, since the PR body deliberately explains the ``connection`` naming.

3. **Trivial:** ``migrations.py:9`` prose still references ``db._SCHEMA``, which no longer resolves at that path (it's ``db.schema._SCHEMA`` now). Doc-only.

**Bottom line:** safe to merge. The only thing I'd genuinely consider is restoring the security-rationale comments (#1), and that's fine as a later cleanup.

---

Want me to **post this on #82** (as an LGTM-with-notes)? And still open from before: should I **keep polling muneem for new PRs** ? if so, what cadence, and auto-post or surface here first?

### 92. user (2026-06-25T17:11:47.314Z)

please post it as a comment before merging and yes, pls poll on other PRs for the same repo and do a review for them

### 93. assistant (2026-06-25T17:12:39.769Z)

I'll post the #82 review now (before you merge), then check for any other open PRs and set up the recurring poll.

### 94. assistant (2026-06-25T17:12:51.182Z)

## LGTM ? ? faithful split, behavior provably preserved (non-blocking notes below)

Since this moves crypto/key-handling code, I verified it at the AST level rather than taking "it's a refactor" on faith ? diffed every symbol in the pre-split ``db.py`` against the merged ``db/`` modules. **After stripping docstrings, all 33 symbols are executable-identical** except two, both verified by hand:

- **The 4 ``list_`` per-bank functions** now share a new ``_JOIN_STATUS`` constant. The f-string reconstructs **byte-identical SQL** ? same ``JOIN documents d ON d.id = t.document_id ORDER BY t.id DESC LIMIT ?``, same ``(limit,)`` binding. No result change.

- **``_SCHEMA``** differs by exactly one thing: a 2-line ``--`` SQL comment was removed. SQLite ignores ``--`` comments, so **the DDL is semantically identical** (same tables, columns, constraints).

Everything else ? ``encrypt_database``, ``rekey_database``, ``backup_database``, ``connect``, ``connect_encrypted``, ``_open_encrypted``, key validation, all CRUD ? is byte-faithful (only docstrings/comments changed).

**Public surface intact:** ``db/___init___.py`` re-exports all 26 public names **plus** ``EncryptedDatabaseError``, so ``from muneem import db as dbmod; dbmod.X`` keeps working ? including the ``dbmod.EncryptedDatabaseError`` caught in ``cli/db.py`` and ``cli/_shared.py``. Private symbols stay in submodules, and the **two test edits** are exactly the right minimal follow-through (patch ``_open_encrypted`` at ``muneem.db.connection``, import ``_SCHEMA`` from ``muneem.db.schema``) ? no test weakened. No non-test code references a now-private internal.

**CI fully green:** Ubuntu 3.11/3.12/3.13 + Windows 3.13 (the Windows job exercises the ``sqlcipher3`` wheel + crypto tests) ? ruff / format / mypy / 478 tests / cov ?80%.

## Non-blocking notes (fine as a follow-up)



1. **This is described as a pure split but also bundles editorial trims** ? several **security-rationale comments** were dropped from the crypto/connection code: ``_raw_key_pragma`'s` **"MUST be the first statement on the connection"**, ``_open_encrypted`'s` **"forces key verification"**, ``encrypt_database`'s` **destructive-step-ordering note**, ``connect_encrypted`'s` **rekey self-healing rationale**, and the **"mint a key only for a brand-new ledger"** warning. Behavior-neutral, but those documented **"why the code is safe"**. For security-critical code I'd lean toward keeping them ? and keeping a split purely mechanical also makes it trivially reviewable.
2. **Doc nit:** the ``db/___init___.py`` docstring lists the module as ``connect``, but it's actually ``connection.py`` (ironic, since the PR body deliberately explains the ``connection`` naming).
3. **Trivial:** ``migrations.py:9`` prose still references ``db._SCHEMA``, which no longer resolves at that path (now ``db.schema._SCHEMA``). Doc-only.

Safe to merge. The only thing I'd genuinely consider is restoring the security-rationale comments (#1) in a later pass.

#### 95. user (2026-06-25T17:12:51.204Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\pr82-review.md (file state is current in your context ? no need to Read it back)

#### 96. user (2026-06-25T17:13:00.629Z)

<https://github.com/avidullu/muneem/pull/82#issuecomment-4802266039>

#### 97. user (2026-06-25T17:13:02.392Z)

```
=== ALL open PRs on muneem ===
[{"author":{"id":"MDQ6VXNlcjE2NzQ4Nzc=","is_bot":false,"login":"avidullu","name":"Avi"},"createdAt":"2026-06-25T17:09:10Z","headRefName":"docs/archive-roadmap","isDraft":false,"number":84,"title":"docs: archive stale ROADMAP §§ 4-6 (original Track A/B plan)"}, {"author":{"id":"MDQ6VXNlcjE2NzQ4Nzc=","is_bot":false,"login":"avidullu","name":"Avi"},"createdAt":"2026-06-25T17:06:10Z","headRefName":"feat/axis-cc-total-due","isDraft":false,"number":83,"title":"feat: add total-due cross-check to Axis CC parser"}, {"author":{"id":"MDQ6VXNlcjE2NzQ4Nzc=","is_bot":false,"login":"avidullu","name":"Avi"},"createdAt":"2026-06-25T17:02:17Z","headRefName":"refactor/split-db","isDraft":false,"number":82,"title":"refactor: split db.py (~590 lines) into db/subpackage"}]
```

#### 98. assistant (2026-06-25T17:13:15.429Z)

#82 posted: [comment](<https://github.com/avidullu/muneem/pull/82#issuecomment-4802266039>). Two more open PRs found ? **#83** (a **feature**: Axis CC total-due cross-check) and **#84** (docs: archive ROADMAP). #83 is a behavior change, so it gets a full correctness review. Let me gather both.

#### 99. user (2026-06-25T17:13:23.826Z)

```
##### PR #83 #####
{"additions":81,"baseRefName":"main","body":"Adds a weak arithmetic verification to the Axis credit card parser, closing the validation gap identified in the code-health report.\n\n### Problem\nAxis CC statements have no running balance column, so the parser could never verify its output ? all rows persisted as `UNVERIFIED`. There was no arithmetic to catch a misread amount.\n\n### Solution\nExtract the printed `Total Amount Due` from the CC statement (regex over first 3 pages of text), then cross-check: `(debits) ? (credits) ? total_due` (within the standard half-paisa tolerance).\n\n- Match ? parse graduates to `RECONCILED` (trusted)\n- Mismatch ? stays `UNVERIFIED` (rows stored but flagged)\n- No total-due line ? stays `UNVERIFIED` (older format, backward-compatible)\n\n### Changes\n- `_TOTAL_DUE_RE` regex for Indian-currency total-due lines\n- `parse_axis_cc` returns 4-tuple: `(card_number, rows, bad_amounts, total_due)`\n- `parse_and_insert` performs the cross-check and sets confidence accordingly\n- Updated 3 existing tests for the 4-tuple signature\n- Added 2 new tests: match ? reconciled, mismatch ? unverified\n\n### Verification\n- ? ruff/mypy clean\n- ? 480 tests pass (+2 new)","changedFiles":3,"deletions":14,"headRefName":"feat/axis-cc-total-due","mergeable":"MERGEABLE","number":83,"title":"feat: add total-due cross-check to Axis CC parser"}
```

```

--- CI ---
test (ubuntu-latest, 3.11)      pass    37s
https://github.com/avidullu/muneem/actions/runs/28187232776/job/83492203684
test (ubuntu-latest, 3.12)      pass    48s
https://github.com/avidullu/muneem/actions/runs/28187232776/job/83492203719
test (ubuntu-latest, 3.13)      pass    38s
https://github.com/avidullu/muneem/actions/runs/28187232776/job/83492203739
test (windows-latest, 3.13)     pass    1m37s
https://github.com/avidullu/muneem/actions/runs/28187232776/job/83492203761
ocr      skipping                0
https://github.com/avidullu/muneem/actions/runs/28187222630/job/83492192250
ocr      skipping                0
https://github.com/avidullu/muneem/actions/runs/28187232776/job/83492204567
test (ubuntu-latest, 3.11)      pass    37s
https://github.com/avidullu/muneem/actions/runs/28187222630/job/83492168912
test (ubuntu-latest, 3.12)      pass    47s
https://github.com/avidullu/muneem/actions/runs/28187222630/job/83492168899
--- files ---
src/muneem/parsers/axis.py
tests/parsers/test_axis.py
tests/test_provenance.py

```

```

##### PR #84 #####
{"additions":104,"baseRefName":"main","body":"Archives the stale § 4-6 (Track A/B, Milestone C/D
plan) from ROADMAP.md to a new archive file.\n\n### Problem\nROADMAP.md had ~90 lines of
\"original 2026-05 plan, kept as history\" ? the old Track A/B and Milestone C/D build-order
diagram, milestone checkboxes, and generalize backlog. These predate the Phase table in § 0 and
were explicitly marked as historical. Keeping them inline added noise for readers who should be
looking at the live § 0 table.\n\n### Changes\n- **New**: `docs/archive/roadmap-2026-05.md` ?
the original 2026-05 plan preserved for reference\n- **ROADMAP.md**: §§ 4-6 replaced with a
short pointer to the archive\n\nFixes recommendation #8 from code-
health-2026-06-25.md.\",\"changedFiles\":2,\"deletions\":86,\"headRefName\":\"docs/archive-
roadmap\",\"mergeable\":\"MERGEABLE\",\"number\":84,\"title\":\"docs: archive stale ROADMAP §§ 4-6
(original Track A/B plan)\"}

```

```

--- CI ---
test (ubuntu-latest, 3.11)      pass    31s
https://github.com/avidullu/muneem/actions/runs/28187403664/job/83492798700
test (ubuntu-latest, 3.12)      pass    46s
https://github.com/avidullu/muneem/actions/runs/28187403664/job/83492798679
test (ubuntu-latest, 3.13)      pass    47s
https://github.com/avidullu/muneem/actions/runs/28187403664/job/83492798631
test (windows-latest, 3.13)     pass    1m20s
https://github.com/avidullu/muneem/actions/runs/28187403664/job/83492798716
ocr      skipping                0
https://github.com/avidullu/muneem/actions/runs/28187403664/job/83492799716
test (ubuntu-latest, 3.11)      pass    36s
https://github.com/avidullu/muneem/actions/runs/28187391639/job/83492759157
test (ubuntu-latest, 3.12)      pass    54s
https://github.com/avidullu/muneem/actions/runs/28187391639/job/83492759140
test (ubuntu-latest, 3.13)      pass    40s
https://github.com/avidullu/muneem/actions/runs/28187391639/job/83492759126
--- files ---
docs/ROADMAP.md
docs/archive/roadmap-2026-05.md

```

#### 100. assistant (2026-06-25T17:13:39.539Z)

Both are green. **#84** is docs (archive move ? I'll verify content faithfulness). **#83** is a real feature touching money arithmetic in a parser ? that needs a careful correctness review. Let me dig into **#83**'s diff and the parser context, and check **#84**'s move faithfulness in parallel.

#### 101. user (2026-06-25T17:13:41.463Z)

```

=== PR #83 FULL DIFF ===
diff --git a/src/muneem/parsers/axis.py b/src/muneem/parsers/axis.py
index ab9d68b..11455b5 100644
--- a/src/muneem/parsers/axis.py
+++ b/src/muneem/parsers/axis.py
@@ -32,6 +32,14 @@
 
 _CC_DATE_RE = re.compile(r"\d{2}/\d{2}/\d{4}") # Axis credit card: dd/mm/yyyy
 _CARD_RES = [re.compile(r"(\d{4}XXXX\d{4})"), re.compile(r"(\d{6}\*\d{4})")]
+# Axis CC statements print the total amount due in forms like:
+#   Total Amount Due           12,345.67
+#   Total Amount Due (Rs.)     12,345.67
+#   Total Payment Due          12,345.67 Cr
+_TOTAL_DUE_RE = re.compile(
+    r"Total\s+(?:Amount|Payment)\s+Due\s*(?:\s*(Rs\.\s*))?\s*([\d,]+\.\s*\d{0,2})",
+    re.IGNORECASE,
+)

def _first_match(text: str, patterns: list[re.Pattern]) -> str | None:
@@ -160,19 +168,33 @@ def _rows_from_cc_tables(tables: list) -> tuple[list[AxisCCTxn], int]:
    return rows, bad_amounts

-def parse_axis_cc(pdf_path: Path, password: str | None = None) -> tuple[str, list[AxisCCTxn],
int]:
-    """Open an Axis credit-card statement; return (card_number, transactions, bad_amounts)."""
+def parse_axis_cc(
+    pdf_path: Path, password: str | None = None
+) -> tuple[str, list[AxisCCTxn], int, float | None]:
+    """Open an Axis credit-card statement.
+
+    Returns ``(card_number, transactions, bad_amounts, total_due)``.
+    ``total_due`` is the printed "Total Amount Due" from the statement,
+    extracted via regex from the first few pages. It is ``None`` if the
+    regex couldn't find it (older format, or a statement that doesn't use
+    that phrasing). The caller can cross-check it against the sum of
+    extracted transaction amounts for a weak arithmetic verification.
+
+    """
+    card_number = "UNKNOWN"
+    rows: list[AxisCCTxn] = []
+    bad_amounts = 0
+    total_due: float | None = None
+    with pdfplumber.open(pdf_path, password=password) as pdf:
+        head = "\n".join((pdf.pages[i].extract_text() or "") for i in range(min(3,
len(pdf.pages))))
+        card_number = _first_match(head, _CARD_RES) or "UNKNOWN"
+        m = _TOTAL_DUE_RE.search(head)
+        if m:
+            total_due = parse_amount(m.group(1).replace(",", ""))
+        for page in pdf.pages:
+            page_rows, page_bad = _rows_from_cc_tables(page.extract_tables())
+            rows.extend(page_rows)
+            bad_amounts += page_bad
+    return card_number, rows, bad_amounts
+    return card_number, rows, bad_amounts, total_due

@register_parser
@@ -183,16 +205,19 @@ def can_handle(self, text: str, sender: str | None = None, subject: str |
None =
    def parse_and_insert(
        self, pdf_path: Path, conn: sqlite3.Connection, doc_id: int, password: str | None =
None
    ) -> int:

```

```

- """Persist CC rows with an explicit ``unverified`` verdict ? never as silently fine.
+ """Persist CC rows; cross-check against the printed total-due when available.
+
+ A credit-card bill has no running balance, so no full arithmetic oracle can endorse
+ the parse. But if the statement prints a "Total Amount Due", we sum the extracted
+ debits/credits and compare: if the net matches the printed total (within tolerance),
+ the parse graduates to RECONCILED. Without a total-due line (older formats), rows
+ persist as UNVERIFIED ? surfaced by listings but never silently trusted.

- A credit-card bill has no running balance, so no arithmetic oracle can endorse the
- parse (Confidence.UNVERIFIED, surfaced by every listing). What CAN be checked is
- checked: a transactional row whose amount cell does not parse means the layout read
- is suspect ? the parse is quarantined (store nothing, ``unreconciled``) rather than
- persisting a wrong amount the math could never catch.
+ A transactional row whose amount cell does not parse means the layout read is suspect
+ ? the entire parse is quarantined (store nothing, UNRECONCILED).
+ """
+
+ try:
-     card_number, rows, bad_amounts = parse_axis_cc(pdf_path, password)
+     card_number, rows, bad_amounts, total_due = parse_axis_cc(pdf_path, password)
+ except ParseError:
+     raise
+ except Exception as exc:
@@ -208,6 +233,16 @@ def parse_and_insert(
+     parser_version=__version__,
+ )
+     return 0
+
+ # Cross-check against the printed total-due when available.
+ confidence = Confidence.UNVERIFIED
+ if total_due is not None:
+     net = sum(r.amount if r.type == "Dr" else -r.amount for r in rows)
+     from .validation import BALANCE_TOLERANCE
+
+     if abs(net - total_due) <= BALANCE_TOLERANCE:
+         confidence = Confidence.RECONCILED
+
+ for t in rows:
+     dbmod.insert_axis_cc_transaction(
+         conn,
@@ -222,7 +257,7 @@ def parse_and_insert(
+     dbmod.record_parse(
+         conn,
+         doc_id,
-         confidence=Confidence.UNVERIFIED.value,
+         confidence=confidence.value,
+         parser="AxisCreditCardParser",
+         parser_version=__version__,
+     )
diff --git a/tests/parsers/test_axis.py b/tests/parsers/test_axis.py
index 1820e4d..77bfb82 100644
--- a/tests/parsers/test_axis.py
+++ b/tests/parsers/test_axis.py
@@ -142,6 +142,7 @@ def
test_rows_from_cc_tables_counts_unparseable_amounts_instead_of_coercing_to_z

def _cc_parse_and_insert(tmp_path, monkeypatch, parse_result):
+ """parse_result is (card_number, rows, bad_amounts, total_due)."""
+ monkeypatch.setattr(axis, "parse_axis_cc", lambda *a, **k: parse_result)
+ with dbmod.connect(tmp_path / "t.db") as conn:
+     dbmod.init_db(conn)
@@ -155,7 +156,8 @@ def _cc_parse_and_insert(tmp_path, monkeypatch, parse_result):

def test_cc_parse_and_insert_records_explicit_unverified(tmp_path, monkeypatch):

```

```

    txns = [axis.AxisCCTxn("05/02/2026", "Swiggy", "Food", 250.0, "Dr")]
-   n, status, stored, prov = _cc_parse_and_insert(tmp_path, monkeypatch, ("CARD1", txns, 0))
+   # No total-due line ? stays UNVERIFIED (no arithmetic to cross-check)
+   n, status, stored, prov = _cc_parse_and_insert(tmp_path, monkeypatch, ("CARD1", txns, 0,
None))
    assert n == 1 and len(stored) == 1
    # No oracle exists for a CC bill: rows persist, but explicitly flagged ? never trusted.
    assert status == "unverified"
@@ -164,7 +166,37 @@ def test_cc_parse_and_insert_records_explicit_unverified(tmp_path,
monkeypatch):

    def test_cc_parse_and_insert_quarantines_on_unparseable_amounts(tmp_path, monkeypatch):
        txns = [axis.AxisCCTxn("05/02/2026", "Swiggy", "Food", 250.0, "Dr")]
-        n, status, stored, _ = _cc_parse_and_insert(tmp_path, monkeypatch, ("CARD1", txns, 2))
+        n, status, stored, _ = _cc_parse_and_insert(tmp_path, monkeypatch, ("CARD1", txns, 2,
None))
        # A bad amount means the layout read is suspect; with no math to catch a misread,
        # store nothing rather than possibly-wrong amounts.
        assert n == 0 and stored == [] and status == "unreconciled"
+
+
+def test_cc_total_due_cross_check_reconciles_on_match(tmp_path, monkeypatch):
+    """When the printed total-due matches the net of extracted amounts, the parse graduates
+    to RECONCILED ? the weak arithmetic verification for a CC bill (no running balance)."""
+    txns = [
+        axis.AxisCCTxn("01/03/2026", "Swiggy", "Food", 250.0, "Dr"),
+        axis.AxisCCTxn("02/03/2026", "Amazon", "Shopping", 150.0, "Dr"),
+        axis.AxisCCTxn("05/03/2026", "Payment", "Payment", 100.0, "Cr"),
+    ]
+    # Net: 250 + 150 - 100 = 300.0
+    n, status, stored, prov = _cc_parse_and_insert(tmp_path, monkeypatch, ("CARD1", txns, 0,
300.0))
+    assert n == 3 and len(stored) == 3
+    assert status == "reconciled"
+    assert prov["confidence"] == "reconciled"
+
+
+def test_cc_total_due_cross_check_stays_unverified_on_mismatch(tmp_path, monkeypatch):
+    """When the printed total-due does NOT match, the parse stays UNVERIFIED ? the math
+    doesn't add up, so we don't endorse it."""
+    txns = [
+        axis.AxisCCTxn("01/03/2026", "Swiggy", "Food", 250.0, "Dr"),
+    ]
+    n, status, stored, _ = _cc_parse_and_insert(
+        tmp_path,
+        monkeypatch,
+        ("CARD1", txns, 0, 500.0), # printed total ? 250
+    )
+    assert n == 1 and len(stored) == 1
+    assert status == "unverified" # rows stored but not trusted
diff --git a/tests/test_provenance.py b/tests/test_provenance.py
index 2959942..c65138c 100644
--- a/tests/test_provenance.py
+++ b/tests/test_provenance.py
@@ -83,7 +83,7 @@ def test_unreconciled_parse_records_unreconciled_provenance(monkeypatch):

    def test_credit_card_parser_records_explicit_unverified_confidence(monkeypatch):
-        monkeypatch.setattr(axis, "parse_axis_cc", lambda *a, **k: ("CARD1234", [], 0))
+        monkeypatch.setattr(axis, "parse_axis_cc", lambda *a, **k: ("CARD1234", [], 0, None))
        with dbmod.connect(":memory:") as conn:
            dbmod.init_db(conn)
            doc = _doc(conn)

```

#83 is safe by design (it only \*upgrades\* trust on an exact arithmetic match), but it's money-parsing logic, so let me verify my concerns precisely before writing them up: test the regex against its own three documented formats (esp. the "Cr" credit-balance case), check `BALANCE\_TOLERANCE`, and confirm the `validation` import situation. Plus verify #84's archive move is loss-free.

### 103. user (2026-06-25T17:15:01.141Z)

```
=== axis.py top imports + AxisCCTxn ===
12:from __future__ import annotations
14:import re
15:import sqlite3
16:from dataclasses import dataclass
17:from pathlib import Path
19:import pdfplumber
21:from muneem import __version__
22:from muneem import db as dbmod
23:from muneem.errors import ParseError
24:from muneem.redaction import redact_filename
26:from . import register_parser
27:from .base import StatementParser
28:from .dates import normalize_statement_date
29:from .profiles import AXIS_PROFILE
30:from .savings import SavingsParse, parse_savings, parse_statement_summary, txn_line_predicate
31:from .validation import Confidence, parse_amount
```

```
=== AxisCCTxn dataclass def ===
128:class AxisCCTxn:
129-     txn_date: str
130-     particulars: str
131-     merchant_category: str
132-     amount: float
133-     type: str # "Cr" or "Dr"
134-
135-
136-def _rows_from_cc_tables(tables: list) -> tuple[list[AxisCCTxn], int]:
```

```
=== BALANCE_TOLERANCE value ===
39:BALANCE_TOLERANCE = 0.005
91:     tol: float = BALANCE_TOLERANCE,
129:     rows: Sequence[ConceptRow], opening: float, closing: float, tol: float =
BALANCE_TOLERANCE
184:     rows: Sequence[ConceptRow], concepts: StatementConcepts, tol: float = BALANCE_TOLERANCE
215:     candidates: Sequence[Sequence[ConceptRow]], tol: float = BALANCE_TOLERANCE
234:     "BALANCE_TOLERANCE",
```

```
=== is 'validation' imported at axis.py module level? ===
```

```
12:from __future__ import annotations
14:import re
15:import sqlite3
16:from dataclasses import dataclass
17:from pathlib import Path
19:import pdfplumber
21:from muneem import __version__
22:from muneem import db as dbmod
23:from muneem.errors import ParseError
24:from muneem.redaction import redact_filename
```

### 104. user (2026-06-25T17:15:07.733Z)

```
'Total Amount Due          12,345.67' -> 12,345.67 | trailing-Cr present: False
'Total Amount Due (Rs.)    12,345.67' -> 12,345.67 | trailing-Cr present: False
'Total Payment Due        12,345.67 Cr' -> 12,345.67 | trailing-Cr present: True
'Total Amount Due\n12,345.67' -> 12,345.67 | trailing-Cr present: False
'Minimum Amount Due 500.00\nTotal Amount D' -> 12,345.67 | trailing-Cr present: False
```

```

=== #84: does the archive file contain exactly what ROADMAP lost? ===
REMOVED ## 4. Build order at a glance
REMOVED
REMOVED ```
REMOVED Milestone 0 ? Decisions & scaffolding (no sensitive data)
REMOVED ?
REMOVED ?????????????????????????????????????????????????????????????
REMOVED ? ?
REMOVED TRACK A: Promo spine TRACK B: Gmail ingestion (run in parallel,
REMOVED parse ? store ? query fetch ? stage ? catalog fully decoupled)
REMOVED (committed corpus, CI) (real email, local only)
REMOVED ? ?
REMOVED ?????????????????????????????????????
REMOVED ?
REMOVED Milestone C ? HDfC vertical slice (CONVERGENCE)
REMOVED unlock (hybrid passwords) ? HDfC parser ? transactions ? CLI
REMOVED ? completes DESIGN.md § 7
REMOVED ?
REMOVED ?
REMOVED Milestone D ? Harden & evaluate (the pre-generalize checkpoint)
REMOVED ?
REMOVED ?
REMOVED Generalize (more issuers, local-LLM fallback, holdings, items, analytics, UI)
REMOVED ```
REMOVED
REMOVED Tracks A and B are genuinely independent: A needs no OAuth and no secrets; B needs no
REMOVED parsing. As a single developer you can interleave them in any order ? neither blocks the other.
REMOVED **Both must land before Milestone C**, where a fetched-and-unlocked HDfC PDF flows through the
REMOVED parser for the first time.
REMOVED
REMOVED Effort tags below are rough (S / M / L), for relative sequencing only ? not calendar
REMOVED estimates.
ADDED ## 4?6. Original 2026-05 plan ? archived
REMOVED
REMOVED ## 5. Milestones & tasks
REMOVED
REMOVED ### ? Milestone 0 ? Decisions & scaffolding · DONE
REMOVED
REMOVED - [x] **0.1 (S)** Project skeleton: `pyproject.toml` (uv), package dir (name TBD, e.g.
REMOVED `muneem/`), `ruff` + `pytest` configured. *Do not over-materialize the source tree (DESIGN § 9)
REMOVED ? create only what 0.x needs.*
REMOVED - [x] **0.2 (S)** Offline CI: GitHub Actions (or local pre-commit) that installs deps
REMOVED and runs `pytest` against committed data only. A trivial `test_smoke` proves the loop is green.
REMOVED - [x] **0.3 (S)** Config + secrets foundation: a config loader (env + a gitignored
REMOVED `config.toml`), and a `keyring`-backed secret accessor stub. Confirm `.gitignore` covers
REMOVED `config.toml` and any new secret paths.
REMOVED - [x] **0.4 (S)** `gitignore` audit: verify the proposed local-only dirs
REMOVED (`testdata/statements/`, staging dir, `*.db`) are blocked. Add a
REMOVED `testdata/statements/.gitkeep`-style note or README so the expected-but-absent dir is
REMOVED documented.
REMOVED - **Exit:** `uv run pytest` is green in CI on a clean checkout; repo layout agreed.
REMOVED
REMOVED ### ? Track A ? Promo spine: parse ? store ? query · DONE (committed corpus, green in
REMOVED CI)
REMOVED
REMOVED > Resurrects the Email-Mining deals feature as muneem's first working output. Zero
REMOVED secrets, zero OAuth.
REMOVED
REMOVED - [x] **A.1 (M)** Port Email-Mining's extraction into a real module: **pdfminer.six**
REMOVED text + **pypdfium2** image extraction (PyMuPDF is banned/AGPL), size filter, dedup. Proper
REMOVED arguments/config; returns a value instead of one-shot stdout. *(Done ? `muneem.extract`)*
REMOVED - [x] **A.2 (S)** Minimal schema slice + migration v1: `documents` (id, source,
REMOVED external_id, sender, subject, received_date, doc_type, sha256, path, status, created_at) and
REMOVED `offers` (id, document_id, merchant, promo_code, description, expiry_date, redeemed,
REMOVED source_text). Full schema is future work (`docs/schema.md`).

```

REMOVED - [x] \*\*A.3 (M)\*\* Promo extraction, \*\*deterministic-first\*\*: regex/heuristics for promo codes + expiry dates over the unified text. Where deterministic extraction misses, an \*\*LLM fallback\*\* is permitted \*because promo is non-sensitive\* ? but that pulls in a cloud service, so it's gated on a quick okay (reuse Email-Mining's Gemini path, go local via Ollama, or use Anthropic ? your call when we get there).

REMOVED - [x] \*\*A.4 (M)\*\* `muneem ingest <pdf>` writes a `documents` + `offers` row; `muneem offers list` reads them back. \*\*Ground-truth + regression test:\*\* establish the known promo codes for each of the 10 corpus PDFs, then a `pytest` that runs the pipeline over `testdata/promo-emails/` and asserts recovery. (Building the ground-truth answer key is itself part of this task.)

REMOVED - \*\*Exit:\*\* `ingest ? SQLite ? offers list` works end-to-end; green regression test on the committed corpus. Stages 3?4?5 proven with zero risk.

REMOVED

REMOVED ### Track B ? Gmail ingestion · PARTIAL (OAuth + search + download ?; `muneem fetch` B.4 pending)

### 105. assistant (2026-06-25T17:15:36.405Z)

Both verifications are conclusive. For #83 I confirmed empirically: the regex \*\*captures\*\* `12,345.67` from the "...Cr" credit-balance line but silently drops the "Cr"\*\*\* ? so that documented case can never reconcile (safe direction, but the comment implies it's handled). And `validation` is already imported at module top (line 31), so the function-local `BALANCE\_TOLERANCE` import is just inconsistent, not circular-dodging. Let me confirm #84's archive is loss-free before writing both up.

### 106. user (2026-06-25T17:15:42.289Z)

```
=== archive file headers present? (content faithfully moved) ===
1:# muneem ? Original 2026-05 Roadmap (Track A/B, Milestone C/D)
10:## Build order at a glance (original)
13:Milestone 0 ? Decisions & scaffolding (no sensitive data)
23:      Milestone C ? HDFC vertical slice (CONVERGENCE)
28:      Milestone D ? Harden & evaluate (the pre-generalize checkpoint)
34:Tracks A and B are genuinely independent: A needs no OAuth and no secrets; B needs no
parsing. As a single developer you can interleave them in any order ? neither blocks the other.
**Both must land before Milestone C**, where a fetched-and-unlocked HDFC PDF flows through the
parser for the first time.
40:## Milestones & tasks (original)
42:### ? Milestone 0 ? Decisions & scaffolding · DONE
50:### ? Track A ? Promo spine: parse ? store ? query · DONE (committed corpus, green in CI)
60:### Track B ? Gmail ingestion · PARTIAL (OAuth + search + download ?; `muneem fetch` B.4
pending)
71:### Milestone C ? HDFC vertical slice · IN PROGRESS (HDFC/Axis/AU parsers ?, unlock-wiring
C.1 ?; unified schema C.3 pending)
80:### Milestone D ? Harden & evaluate · the pre-generalize checkpoint
```

```
=== archive file line count vs ROADMAP removed (~86) ===
98
```

```
=== ROADMAP.md pointer that replaced the sections ===
133:## 4?6. Original 2026-05 plan ? archived
134-
135-The original Track A/B and Milestone C/D plan (build-order diagram, milestone
136-checkboxes, and generalize backlog) has been **moved to
137:[`docs/archive/roadmap-2026-05.md`](./archive/roadmap-2026-05.md)**. The live
138-plan is **$ 0 (Phases)** above ? it supersedes the old milestone letters, which
139-are unrelated to the current phase letters.
140-
141----
```

```
=== any dangling inbound refs to the removed ROADMAP sections elsewhere in docs? ===
docs/CURRENT_STATE.md:15:The ingestion pipeline supports a generic `SearchCriteria` interface
and multiple providers (Track B plumbing):
docs/feedback-antigravity-response.md:34:  `muneem fetch` (B.4) ships (ROADMAP Track B;
feedbackJune2026 P7).
```



```
docs/ROADMAP.md:14:> **Track A** (promo spine: `ingest` ? SQLite ? `offers list`, green
regression) ? are done.
docs/ROADMAP.md:15:> **Track B:** Gmail OAuth2 + search + attachment download ?; the `muneem
fetch` CLI (B.4) is not
docs/ROADMAP.md:28:> doc now defers here for status). Naming note: the **Track A/B** and
**Milestone C/D** sections
docs/ROADMAP.md:129:| LLM (promo only, non-sensitive) | deterministic-first; cloud LLM **only if
needed and approved** | See § 5 Track A. Sensitive-doc LLM is parked (§ 9 / DESIGN § 8.3). |
docs/ROADMAP.md:146:      Milestone C ? HDFC vertical slice (CONVERGENCE)
docs/ROADMAP.md:151:      Milestone D ? Harden & evaluate (the pre-generalize checkpoint)
docs/ROADMAP.md:157:Tracks A and B are genuinely independent: A needs no OAuth and no secrets; B
needs no parsing. As a single developer you can interleave them in any order ? neither blocks
the other. **Both must land before Milestone C**, where a fetched-and-unlocked HDFC PDF flows
through the parser for the first time.
docs/ROADMAP.md:173:### ? Track A ? Promo spine: parse ? store ? query · DONE (committed corpus,
green in CI)
```

#### 107. assistant (2026-06-25T17:15:59.642Z)

The "dangling refs" hits at `ROADMAP.md:146-173` are from my **working tree** (checked out on a different branch), not PR #84 ? let me confirm against #84's branch directly that the old sections are actually gone (not duplicated), since that's the one real risk in an archive-move.

#### 108. user (2026-06-25T17:16:04.348Z)

```
=== working tree is on branch: ===
refactor/split-db
```

```
=== ON PR #84 BRANCH: does ROADMAP.md still contain the OLD sections (should be GONE)? ===
>>> NONE found in ROADMAP on #84 branch ? sections fully moved, no duplication (good)
```

```
=== ROADMAP.md total lines: base vs #84 branch ===
base:   333
#84:   253
```

#### 109. assistant (2026-06-25T17:17:26.874Z)

All verified. **#84** is a clean, loss-free archive move (ROADMAP 280 lines, 98-line archive with every section header preserved, correct pointer link, no duplication). **#83** is a safe feature with four worth-noting items, two of them real (the dropped "Cr" sign + the untested regex). Let me write and post both reviews.

#### 110. assistant (2026-06-25T17:17:31.452Z)

### LGTM ? ? faithful, loss-free archive move

Verified the content was *moved*, not lost or duplicated:

- **No content lost:** `docs/archive/roadmap-2026-05.md` (98 lines) carries every section that left ROADMAP ? build-order diagram, Milestone 0, Track A, Track B, Milestones & tasks, Milestone C/D, generalize backlog. All headers present.
- **No duplication:** on this branch, ROADMAP.md no longer contains any of the moved sections (grep for "Build order at a glance", "Milestone C ? HDFC", "Track A ? Promo spine", "Milestones & tasks" ? none). ROADMAP went 333 ? 253 lines (?80).
- **Pointer is correct:** the replacement `## 4?6` section links to `./archive/roadmap-2026-05.md` (valid relative path) and correctly redirects readers to the live § 0 Phase table.
- The brief Track A/B status mentions kept near § 0 are the live summary, not the archived detail ? correct to keep.

Docs-only, CI green. Good to merge.

#### 111. user (2026-06-25T17:17:31.477Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\pr84-review.md (file state is current in your context ? no need to Read it back)

## 112. assistant (2026-06-25T17:17:47.101Z)

### Review ? feature is sound and safe; a few things worth tightening

The design is the right shape: the cross-check only \*upgrades\* a CC parse to `RECONCILED` when `?(Dr) ? ?(Cr)` matches the printed total-due within half a paisa; a mismatch or a missing total-due line stays `UNVERIFIED`, and unparseable amounts still quarantine the whole parse. So there's **no false-trust risk** ? the worst case is under-trusting a correct parse. Float arithmetic is fine here (tolerance `0.005` ? accumulation error) and consistent with the savings parsers. CI is green (480 tests, ruff/mypy) on Ubuntu 3.11/3.12/3.13 + Windows.

Four things to consider ? none are blockers, but the first two are worth doing:

#### 1. The "Cr" (credit-balance) case is documented but not actually handled

The regex comment lists `Total Payment Due 12,345.67 Cr` as a supported form, but the trailing `Cr` is dropped ? I verified `\_TOTAL\_DUE\_RE` captures `12,345.67` (positive) and ignores the `Cr`. So when the customer has a **credit balance** (overpaid; true due is negative), `total\_due` is stored as `+12345.67` while the actual net is ? `12345.67` ? the cross-check can **never** match ? the statement silently stays `UNVERIFIED`. Safe direction, but the comment implies it's handled. Either negate `total\_due` when a trailing `Cr`/`CR` is present, or drop that example from the comment so it doesn't over-promise.

#### 2. `\_TOTAL\_DUE\_RE` has no direct test

All five tests monkeypatch `parse\_axis\_cc` to return a fixed 4-tuple, so they exercise the cross-check arithmetic in `parse\_and\_insert` ? but the **regex itself is never asserted**, and it's the fragile part (three formats, the `Cr` case, number-on-next-line, "Minimum Amount Due" vs "Total Amount Due" disambiguation). A handful of direct `\_TOTAL\_DUE\_RE` cases would lock in the documented formats ? and would have surfaced #1.

#### 3. Worth a one-line note: the check assumes no carried-over balance

`?(Dr) ? ?(Cr)` equals "Total Amount Due" only when there's **no uncaptured previous/opening balance**. A revolving balance (carried from last cycle) makes the net legitimately differ from the printed total, so a correctly-parsed statement stays `UNVERIFIED`. That's fine by design, but a comment naming the limitation would stop a future reader from "fixing" it as a bug.

#### 4. Minor: hoist the `BALANCE\_TOLERANCE` import

`from .validation import BALANCE\_TOLERANCE` is imported inside `parse\_and\_insert`, but `validation` is already imported at module top (line 31: `from .validation import Confidence, parse\_amount`). No circular-import reason ? just add `BALANCE\_TOLERANCE` to that existing line for consistency.

Safe to merge as-is; I'd fold in #1 + #2 either here or as a quick follow-up.

## 113. user (2026-06-25T17:17:47.122Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\pr83-review.md (file state is current in your context ? no need to Read it back)

## 114. user (2026-06-25T17:17:57.451Z)

<https://github.com/avidullu/muneem/pull/84#issuecomment-4802303723>

<https://github.com/avidullu/muneem/pull/83#issuecomment-4802303902>

## 115. assistant (2026-06-25T17:18:05.758Z)

Both posted ([#84](https://github.com/avidullu/muneem/pull/84#issuecomment-4802303723), [#83](https://github.com/avidullu/muneem/pull/83#issuecomment-4802303902)). Now let me set up the recurring poll ? seed a state file of what I've already reviewed (by PR number + head commit, so re-pushes get re-reviewed) so the poll only acts on genuinely new work.

## 116. user (2026-06-25T17:18:10.238Z)

```
=== current open PRs with head SHAs ===
[{"headRef0id":"b5f11d3fd36ec9e7f2d4898c6473945d9abd7aad","number":84,"title":"docs: archive
stale ROADMAP §§ 4-6 (original Track A/B
plan)"},"{"headRef0id":"75d220ca73580bc9b3ade1cc254986960ef7de43","number":83,"title":"feat: add
total-due cross-check to Axis CC
parser"},"{"headRef0id":"8903ec75c46f2e6f195abebad95d70ceb3c56f86","number":82,"title":"refactor:
split db.py (~590 lines) into db/ subpackage"}]
```

```
=== seeded state file (C:/Users/avidu/AppData/Local/Temp/muneem_review_state.json) ===
{"82":"8903ec75c46f2e6f195abebad95d70ceb3c56f86","83":"75d220ca73580bc9b3ade1cc254986960ef7de43"
,"84":"b5f11d3fd36ec9e7f2d4898c6473945d9abd7aad"}
```

## 117. assistant (2026-06-25T17:18:33.703Z)

State seeded. Now I'll schedule the recurring poll.

## 118. user (2026-06-25T17:18:47.108Z)

Next wakeup scheduled for 23:09:00 (in 1213s). Nothing more to do this turn ? the harness re-  
invokes you when the wakeup fires or a task-notification arrives.

## 119. assistant (2026-06-25T17:19:04.335Z)

All three open PRs reviewed and posted, and the poll is live.

### Reviews posted

```
| PR | Type | Verdict | Posted |
|----|-----|-----|-----|
| [#82](https://github.com/avidullu/muneem/pull/82) | refactor (db split) | **LGTM** ? AST-
verified faithful; 3 non-blocking notes (lost security-rationale comments,
`connect`/`connection` docstring nit) |
[comment](https://github.com/avidullu/muneem/pull/82#issuecomment-4802266039) |
| [#83](https://github.com/avidullu/muneem/pull/83) | feature (CC total-due check) | **Safe +
green, not blanket-LGTM** ? 4 notes; the two real ones: dropped "Cr" sign + untested regex |
[comment](https://github.com/avidullu/muneem/pull/83#issuecomment-4802303902) |
| [#84](https://github.com/avidullu/muneem/pull/84) | docs (archive ROADMAP) | **LGTM** ? loss-
free move, no duplication, valid link |
[comment](https://github.com/avidullu/muneem/pull/84#issuecomment-4802303723) |
```

The one I'd actually act on before/after merging is **#83**: the regex comment advertises a  
"Total Payment Due ? Cr" form, but the trailing `Cr` is dropped ? I confirmed empirically it  
captures `12,345.67` as positive, so genuine credit-balance statements can never reconcile (safe  
direction, but the comment over-promises). And the regex itself has zero direct tests. Both are  
easy follow-ups; nothing is unsafe.

### Polling is set up

- **Cadence**: every ~20 min, I re-check muneem for **open PRs** that are new or have new  
commits since I last reviewed them (tracked by PR number + head SHA in a state file, so a re-  
push gets re-reviewed).
- **What it does**: reviews by type (refactor ? AST-faithfulness; feature ? correctness, defer  
merge to you; docs ? faithfulness), **auto-posts** the review comment, and reschedules itself.
- **Guards**: it won't sign off while CI is pending/red, won't touch the shared checkout (read-  
only via `git show`), and stays quiet when there's nothing new.
- Next check fires ~23:09.

A few things worth confirming:

- **Auto-post is on**: The poll posts reviews without checking with you first (that's how I read  
"do a review for them"). Say the word if you'd rather it **surface here first** for your go-  
ahead.
- You can still just tell me "#85 is up" anytime to trigger an immediate review instead of  
waiting for the tick.

- To stop or change cadence, just say so. Note the poll runs while this session is alive ? if you want it to survive restarts, I can set it up as a scheduled cloud agent instead.

#### 120. user (2026-06-25T17:19:06.271Z)

pls poll for the comments on the PRs as well. #82 has responses

#### 121. assistant (2026-06-25T17:19:11.916Z)

Let me read #82's responses first, then wire comment-polling into the loop.

#### 122. user (2026-06-25T17:19:15.654Z)

```
=== #82 issue comments (chronological) ===
--- avidullu @ 2026-06-25T17:13:00Z ---
```

#### LGTM ? ? faithful split, behavior provably preserved (non-blocking notes below)

Since this moves crypto/key-handling code, I verified it at the AST level rather than taking "it's a refactor" on faith ? diffed every symbol in the pre-split `db.py` against the merged `db/` modules. **\*\*After stripping docstrings, all 33 symbols are executable-identical\*\*** except two, both verified by hand:

- **\*\*The 4 `list\_` per-bank functions\*\*** now share a new `\_JOIN\_STATUS` constant. The f-string reconstructs **\*\*byte-identical SQL\*\*** ? same `JOIN` documents `d ON d.id = t.document_id ORDER BY t.id DESC LIMIT ?`, same `(limit,)` binding. No result change.
- **\*\*`\_SCHEMA`\*\*** differs by exactly one thing: a 2-line `--` SQL comment was removed. SQLite ignores `--` comments, so **\*\*the DDL is semantically identical\*\*** (same tables, columns, constraints).

Everything else ? `encrypt\_database`, `rekey\_database`, `backup\_database`, `connect`, `connect\_encrypted`, `\_open\_encrypted`, key validation, all CRUD ? is byte-faithful (only docstrings/comments changed).

**\*\*Public surface intact:\*\*** `db/\_\_init\_\_.py` re-exports all 26 public names **\*\*plus** `EncryptedDatabaseError`**\*\***, so `from muneem import db as dbmod; dbmod.X` keeps working ? including the `dbmod.EncryptedDatabaseError` caught in `cli/db.py` and `cli/_shared.py`. Private symbols stay in submodules, and the **two test edits are exactly the right minimal follow-through** (patch `_open_encrypted` at `muneem.db.connection`, import `_SCHEMA` from `muneem.db.schema`) ? no test weakened. No non-test code references a now-private internal.`

**\*\*CI fully green:\*\*** Ubuntu 3.11/3.12/3.13 + Windows 3.13 (the Windows job exercises the sqlcipher3 wheel + crypto tests) ? ruff / format / mypy / 478 tests / cov ?80%.

#### Non-blocking notes (fine as a follow-up)

1. **\*\*This is described as a pure split but also bundles editorial trims\*\*** ? several **\*\*security-rationale comments\*\*** were dropped from the crypto/connection code: `\_raw\_key\_pragma`'s **\*\*"MUST be the first statement on the connection"**, `\_open\_encrypted`'s **\*\*"forces key verification"**, `encrypt\_database`'s destructive-step-ordering note, `connect\_encrypted`'s rekey self-healing rationale, and the "mint a key only for a brand-new ledger" warning. Behavior-neutral, but those documented **\*\*why the code is safe\*\***. For security-critical code I'd lean toward keeping them ? and keeping a split purely mechanical also makes it trivially reviewable.
2. **\*\*Doc nit:\*\*** the `db/\_\_init\_\_.py` docstring lists the module as `connect`, but it's actually `connection.py` (ironic, since the PR body deliberately explains the `connection` naming).
3. **\*\*Trivial:\*\*** `migrations.py:9` prose still references `db.\_SCHEMA`, which no longer resolves at that path (now `db.schema.\_SCHEMA`). Doc-only.

Safe to merge. The only thing I'd genuinely consider is restoring the security-rationale comments (#1) in a later pass.

```
=== #82 review threads (inline) ===
```

### 123. assistant (2026-06-25T17:19:28.628Z)

Interesting ? the only comment I can see is my own review (gh is authed as `avidullu`, same as you). The responses must be in a different channel (a formal PR review, or inline line-comments, or they landed just now). Let me check all of them.

### 124. user (2026-06-25T17:19:34.441Z)

```
=== formal PR reviews on #82 ===
=== ALL issue comments (count + bodies, freshly) ===
--- avidullu @ 2026-06-25T17:13:00Z ---
```

### LGTM ?? faithful split, behavior provably preserved (non-blocking notes below)

Since this moves crypto/key-handling code, I verified it at the AST level rather than taking "it's a refactor" on faith ? diffed every symbol in the pre-split `db.py` against the merged `db/` modules. **\*\*After stripping docstrings, all 33 symbols are executable-identical\*\*** except two, both verified by hand:

- **\*\*The**

```
=== inline review comments (raw count) ===
0
```

### 125. user (2026-06-25T17:19:59.661Z)

```
=== #82 state (merged?) ===
{"closed":false,"mergedAt":null,"nComments":1,"state":"OPEN"}
```

```
=== #82 timeline events (comments/reviews/edits) ===
committed | null | null |
commented | avidullu | 2026-06-25T17:13:00Z | ## LGTM ?? faithful split, behavior provably
preserved (non-blocking notes below)
```

Since this moves crypto/key-handling  
committed | null | null |

```
=== comment counts on #83 and #84 ===
#83 issue-comments: 1  inline: 0
#84 issue-comments: 1  inline: 0
```

### 126. assistant (2026-06-25T17:20:58.322Z)

Let me confirm who `gh` is and whether #82 changed in any other way, then reseed state for comment-tracking.

### 127. user (2026-06-25T17:25:23.900Z)

```
gh identity: avidullu
#82 current head: 8903ec75c46f2e6f195abebad95d70ceb3c56f86 (seeded:
8903ec75c46f2e6f195abebad95d70ceb3c56f86)
```

```
=== reseed state v2: {head, last_comment_ts} per open PR ===
{
  "84": {
    "head": "b5f11d3fd36ec9e7f2d4898c6473945d9abd7aad",
    "last_comment_ts": "2026-06-25T17:17:55Z"
  },
  "83": {
    "head": "75d220ca73580bc9b3ade1cc254986960ef7de43",
    "last_comment_ts": "2026-06-25T17:17:57Z"
  },
  "82": {
    "head": "8903ec75c46f2e6f195abebad95d70ceb3c56f86",
    "last_comment_ts": "2026-06-25T17:20:09Z"
  }
}
```

```
}  
}
```

#### 128. user (2026-06-25T17:30:33.833Z)

Continue from where you left off.

#### 129. assistant (2026-06-25T17:30:33.833Z)

No response requested.

#### 130. user (2026-06-25T17:30:38.671Z)

can you complete the review of the PRs?

#### 131. assistant (2026-06-25T17:30:57.262Z)

Let me do a fresh comprehensive sweep of all open muneem PRs ? new ones, new commits, and any responses on the ones I've reviewed.

#### 132. user (2026-06-25T17:31:13.529Z)

=== saved state ===

```
{  
  "84": {  
    "head": "b5f11d3fd36ec9e7f2d4898c6473945d9abd7aad",  
    "last_comment_ts": "2026-06-25T17:17:55Z"  
  },  
  "83": {  
    "head": "75d220ca73580bc9b3ade1cc254986960ef7de43",  
    "last_comment_ts": "2026-06-25T17:17:57Z"  
  },  
  "82": {  
    "head": "8903ec75c46f2e6f195abebad95d70ceb3c56f86",  
    "last_comment_ts": "2026-06-25T17:20:09Z"  
  }  
}
```

=== open PRs now (number, head, title, updatedAt) ===

```
#82 draft=false head=8903ec75 upd=2026-06-25T17:20:09Z refactor: split db.py (~590 lines) into  
db/ subpackage  
#83 draft=false head=75d220ca upd=2026-06-25T17:17:57Z feat: add total-due cross-check to Axis  
CC parser  
#84 draft=false head=b5f11d3f upd=2026-06-25T17:17:55Z docs: archive stale ROADMAP §§ 4-6  
(original Track A/B plan)
```

=== per-PR comment/review/inline counts + latest non-bot timestamps ===

```
#84 issue-comments=1 reviews=0 inline=0  
#83 issue-comments=1 reviews=0 inline=0  
#82 issue-comments=2 reviews=0 inline=0
```

#### 133. assistant (2026-06-25T17:31:18.056Z)

#82 now shows \*\*2 comments\*\* (was 1) ? there's the response. Let me read it.

#### 134. user (2026-06-25T17:31:20.147Z)

===== avidullu @ 2026-06-25T17:13:00Z (id 4802266039) =====

#### **LGTM ? ? faithful split, behavior provably preserved (non-blocking notes below)**

Since this moves crypto/key-handling code, I verified it at the AST level rather than taking "it's a refactor" on faith ? diffed every symbol in the pre-split `db.py` against the merged `db/` modules. \*\*After stripping docstrings, all 33 symbols are executable-identical\*\* except

two, both verified by hand:

```
- **The 4 `list_` per-bank functions** now share a new `_JOIN_STATUS` constant. The f-string reconstructs **byte-identical SQL** ? same `JOIN` documents d ON d.id = t.document_id ORDER BY t.id DESC LIMIT ?, same `(limit,)` binding. No result change.
- **`_SCHEMA`** differs by exactly one thing: a 2-line `--` SQL comment was removed. SQLite ignores `--` comments, so **the DDL is semantically identical** (same tables, columns, constraints).
```

Everything else ? `encrypt\_database`, `rekey\_database`, `backup\_database`, `connect`, `connect\_encrypted`, `\_open\_encrypted`, key validation, all CRUD ? is byte-faithful (only docstrings/comments changed).

```
**Public surface intact:** `db/__init__.py` re-exports all 26 public names **plus `EncryptedDatabaseError`**, so `from muneem import db as dbmod; dbmod.X` keeps working ? including the `dbmod.EncryptedDatabaseError` caught in `cli/db.py` and `cli/_shared.py`. Private symbols stay in submodules, and the **two test edits are exactly the right minimal follow-through** (patch `_open_encrypted` at `muneem.db.connection`, import `_SCHEMA` from `muneem.db.schema`) ? no test weakened. No non-test code references a now-private internal.
```

**\*\*CI fully green:\*\*** Ubuntu 3.11/3.12/3.13 + Windows 3.13 (the Windows job exercises the sqlcipher3 wheel + crypto tests) ? ruff / format / mypy / 478 tests / cov ?80%.

### Non-blocking notes (fine as a follow-up)

1. **\*\*This is described as a pure split but also bundles editorial trims\*\*** ? several **\*\*security-rationale comments\*\*** were dropped from the crypto/connection code: `\_raw\_key\_pragma`'s **\*\*"MUST be the first statement on the connection"**, `\_open\_encrypted`'s **\*\*"forces key verification"**, `encrypt\_database`'s destructive-step-ordering note, `connect\_encrypted`'s rekey self-healing rationale, and the "mint a key only for a brand-new ledger" warning. Behavior-neutral, but those documented **\*\*why the code is safe\*\***. For security-critical code I'd lean toward keeping them ? and keeping a split purely mechanical also makes it trivially reviewable.
2. **\*\*Doc nit:\*\*** the `db/\_\_init\_\_.py` docstring lists the module as `connect`, but it's actually `connection.py` (ironic, since the PR body deliberately explains the `connection` naming).
3. **\*\*Trivial:\*\*** `migrations.py:9` prose still references `db.\_SCHEMA`, which no longer resolves at that path (now `db.schema.\_SCHEMA`). Doc-only.

Safe to merge. The only thing I'd genuinely consider is restoring the security-rationale comments (#1) in a later pass.

==== avidullu @ 2026-06-25T17:20:09Z (id 4802320302) ====  
All three review notes addressed in 8903ec7:

1. **\*\*Security-rationale comments restored\*\*** ? `\_validate\_raw\_key` (raw-key?KDF switch warning), `is\_plaintext\_sqlite` (refuse-to-mistreat rationale), `\_open\_encrypted` (no-downgrade guarantee), inline **MUST-be-first/forces-verification**, **connect\_encrypted** (mint-key-only-for-new-ledger), and **connect** (full encrypted-at-rest docstring).
2. **\*\*`\\_\_init\_\_\\_\_\_.py` doc nit\*\*** ? module listing fixed: **connect** ? **connection**.
3. **\*\*migrations.py prose\*\*** ? **db.\_SCHEMA** ? **db.schema.\_SCHEMA**.

ruff/mypy clean, 478 tests pass. Ready for you.

### 135. assistant (2026-06-25T17:32:16.306Z)

The author responded on #82 (17:20): says all three notes were addressed in commit `8903ec7` and it's "ready for you." But that's the **\*\*same short SHA my review was based on\*\*** ? so I need to verify against the live tree whether the fixes are actually present (a force-push could have reused my view of the head). Let me check the three fixes directly and re-confirm behavior is still preserved.

### 136. user (2026-06-25T17:32:22.575Z)

=== current head of refactor/split-db ===

8903ec75c46f2e6f195abebad95d70ceb3c56f86

=== commit log (did a fix commit land?) ===

8903ec7 fix: address PR #82 review ? restore security comments, fix docs

fca3bd7 refactor: split db.py (~590 lines) into db/ subpackage

ec12b59 fix: address PR #80 post-merge review comments (#81)

=== FIX 2: \_\_init\_\_.py docstring ? connect vs connection? ===

6:- ``connection`` ? key helpers, plaintext/encrypted connections

10:All public names are re-exported here, so ``from muneem.db import connect\_encrypted``

14:from muneem.db.connection import ( # noqa: F401

15: connect,

16: connect\_encrypted,

=== FIX 3: migrations.py:9 ? db.\_SCHEMA vs db.schema.\_SCHEMA? ===

- **Baseline.** The current schema (``db.schema.\_SCHEMA``, version ``db.SCHEMA\_VERSION``)

=== FIX 1: security-rationale comments restored in connection.py? ===

22: Anything else would either silently switch SQLCipher from raw-key mode to passphrase-KDF

35: """The SQLCipher raw-key PRAGMA: use the 256-bit key (64 hex chars) directly, no KDF."""

76: conn.execute(\_raw\_key\_pragma(key)) # MUST be the first statement on the connection

78: conn.execute("SELECT count(\*) FROM sqlite\_master") # forces key verification

### 137. assistant (2026-06-25T17:32:49.349Z)

Mystery solved ? the history shows a **new commit landed**: `8903ec7` "fix: address PR #82 review ? restore security comments, fix docs" on top of the `fca3bd7` split I originally reviewed. My state seed captured the post-fix head, so the poll never flagged it ? I owe a review of that fix commit. All three fixes are visibly present. Now the key check: confirm the fix commit changed **only** comments/docstrings/prose and didn't perturb any executable logic.

### 138. user (2026-06-25T17:33:00.225Z)

=== fix commit touched which files? ===

src/muneem/db/\_\_init\_\_.py | 2 +-

src/muneem/db/connection.py | 33 ++++++

src/muneem/migrations.py | 3 +-

3 files changed, 29 insertions(+), 9 deletions(-)

=== executable-AST diff: fca3bd7 db/ vs 8903ec7 db/ (docstrings stripped) ? should be EMPTY ===

executable symbols differing between the two commits: NONE ? fix commit is comment/docstring-only

=== CI on current head 8903ec7 ===

ocr skipping 0

<https://github.com/avidullu/muneem/actions/runs/28187773269/job/83494110217>

test (ubuntu-latest, 3.11) pass 33s

<https://github.com/avidullu/muneem/actions/runs/28187770831/job/83494075453>

test (ubuntu-latest, 3.11) pass 35s

<https://github.com/avidullu/muneem/actions/runs/28187773269/job/83494084741>

test (ubuntu-latest, 3.12) pass 52s

<https://github.com/avidullu/muneem/actions/runs/28187770831/job/83494075439>

test (ubuntu-latest, 3.12) pass 50s

<https://github.com/avidullu/muneem/actions/runs/28187773269/job/83494084765>

test (ubuntu-latest, 3.13) pass 43s

<https://github.com/avidullu/muneem/actions/runs/28187770831/job/83494075536>

### 139. assistant (2026-06-25T17:33:23.536Z)

Fully verified: the fix commit `8903ec7` touched only `\_\_init\_\_.py`, `connection.py`, and `migrations.py`, and the executable-AST diff against the original split is **NONE** ? it's comment/docstring-only, zero logic change. By transitivity with my original review, the PR at current head is still provably behavior-faithful to the pre-split `db.py`. Let me confirm



Windows CI on this head, then post the verification.

#### 140. user (2026-06-25T17:33:27.763Z)

```
=== Windows CI on #82 current head ===
test (windows-latest, 3.13)    pass    1m55s
https://github.com/avidullu/muneem/actions/runs/28187770831/job/83494075449
test (windows-latest, 3.13)    pass    1m20s
https://github.com/avidullu/muneem/actions/runs/28187773269/job/83494084758

=== sanity: all 3 restored security-comment anchors present in connection.py ===
7
(count of restored rationale anchors)
```

#### 141. assistant (2026-06-25T17:33:40.809Z)

Everything checks out ? Windows CI green, 7 rationale anchors restored. Let me post the verification on #82 and update tracking.

#### 142. assistant (2026-06-25T17:33:49.118Z)

### Re-reviewed `8903ec7` ? all three notes verified addressed ? clean LGTM

Checked the fix commit against the split it sits on (`fca3bd7`):

- **\*\*1. Security-rationale comments restored\*\*** ? confirmed in `connection.py`: the `\_validate\_raw\_key` raw-key?passphrase-KDF warning, `is\_plaintext\_sqlite`'s refuse-to-mistreat rationale, `\_open\_encrypted`'s no-downgrade guarantee, the inline `# MUST be the first statement on the connection` + `# forces key verification`, `connect\_encrypted`'s mint-key-only-for-a-new-ledger note, and `connect`'s full encrypted-at-rest docstring. ?
- **\*\*2. `\_\_init\_\_.py` docstring\*\*** ? now lists the module as `connection` (not `connect`). ?
- **\*\*3. `migrations.py:9`\*\*** ? now `db.schema.\_SCHEMA` (was `db.\_SCHEMA`). ?

**\*\*Behavior still provably preserved.\*\*** The fix commit touches only `\_\_init\_\_.py` (1 line), `connection.py` (comments/docstrings), and `migrations.py` (1 line of prose). I re-ran the docstring-stripped AST comparison between `fca3bd7` and `8903ec7`: **\*\*zero executable symbols differ\*\*** ? it's a pure comment/docstring restoration. Combined with the original review (which proved `fca3bd7` is executable-identical to the pre-split `db.py`), the PR at `8903ec7` remains byte-faithful in behavior.

**\*\*CI green on `8903ec7`\*\*** ? Ubuntu 3.11/3.12/3.13 + Windows 3.13 (the Windows job re-exercised the sqlcipher3 wheel + crypto tests).

Nothing left on my side ? good to merge.

<!-- claude-review-bot -->

#### 143. user (2026-06-25T17:33:49.147Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\pr82-followup.md (file state is current in your context ? no need to Read it back)

#### 144. user (2026-06-25T17:34:00.993Z)

<https://github.com/avidullu/muneem/pull/82#issuecomment-4802438806>

```
=== update state: #82 reviewed-head + advance last_comment_ts past my follow-up ===
{
  "84": {
    "head": "b5f11d3fd36ec9e7f2d4898c6473945d9abd7aad",
    "last_comment_ts": "2026-06-25T17:17:55Z"
  },
  "83": {
    "head": "75d220ca73580bc9b3ade1cc254986960ef7de43",
```

```

    "last_comment_ts": "2026-06-25T17:17:57Z"
  },
  "82": {
    "head": "8903ec75c46f2e6f195abebad95d70ceb3c56f86",
    "last_comment_ts": "2026-06-25T17:33:58Z"
  }
}

```

#### 145. assistant (2026-06-25T17:35:09.460Z)

Posted: [#82 follow-up](https://github.com/avidullu/muneem/pull/82#issuecomment-4802438806). Now let me upgrade the recurring poll to also watch for author responses (your explicit ask), since the previous version only watched for new PRs.

#### 146. user (2026-06-25T17:35:39.538Z)

Next wakeup scheduled for 23:26:00 (in 1220s). Nothing more to do this turn ? the harness re-invokes you when the wakeup fires or a task-notification arrives.

#### 147. assistant (2026-06-25T17:35:59.846Z)

Done ? and the "#82 has responses" is fully resolved.

### What the responses were

The author pushed a **new commit** `8903ec7`` ("fix: address PR #82 review") on top of the split I originally reviewed (`fca3bd7``) and commented that all three notes were addressed. My poll missed it because I'd seeded state with the post-fix head ? exactly the gap you flagged by asking me to watch comments too.

### Re-review of the fix commit (posted: [#82 follow-up](https://github.com/avidullu/muneem/pull/82#issuecomment-4802438806))

Verified all three, and that the fix is behavior-safe:

- ? **Security-rationale comments restored** (7 anchors back in `connection.py`` ? KDF-switch warning, MUST-be-first, forces-verification, no-downgrade, mint-key-only, encrypted-at-rest docstring)
- ? **\_\_init\_\_.py` docstring** ? `connection`` (not `connect``)
- ? **migrations.py:9`** ? `db.schema._SCHEMA``
- ? **Zero executable change** ? docstring-stripped AST diff between `fca3bd7`` and `8903ec7`` is empty; it's pure comment/docstring restoration, so the PR stays byte-faithful to the pre-split `db.py``
- ? **CI green on 8903ec7`** incl. Windows

? Clean LGTM, ready to merge. Nothing left on my side.

### Poll upgraded to watch comments

The recurring poll (every ~20 min) now does **both**:

1. **New/updated PRs** ? review by type.
2. **Author responses on reviewed PRs** ? engage: if a reply says "addressed in `<commit>``", I re-review that commit (diff + AST-faithfulness/finding-recheck + CI) and confirm; if it's a question or pushback, I answer it.

To handle the fact that `gh`` and you are the **same** `avidullu`` account, I tag every comment I post with a hidden `<!-- claude-review-bot -->`` marker ? the poll treats any newer, unmarked comment as your message to respond to. State tracks `{head, last_comment_ts}`` per PR so re-pushes and replies both get caught. Next tick ~23:26.

**Current status:** #82 fully addressed (LGTM); #83 (feature ? safe, my 2 substantive notes still open for you) and #84 (docs ? LGTM) are reviewed and green, no responses yet. Want me to keep #83's notes as-is, or would you like a quick follow-up PR for the "Cr"-sign + regex-test items?

#### 148. user (2026-06-25T17:39:18.958Z)

can you review the comments on #83

149. user (2026-06-25T17:39:31.837Z)

```
=== #83 issue comments (chronological) ===  
===== avidullu @ 2026-06-25T17:17:57Z (id 4802303902) =====
```

## Review ? feature is sound and safe; a few things worth tightening

The design is the right shape: the cross-check only \*upgrades\* a CC parse to `RECONCILED` when `?(Dr) ??(Cr)` matches the printed total-due within half a paisa; a mismatch or a missing total-due line stays `UNVERIFIED`, and unparseable amounts still quarantine the whole parse. So there's **\*no false-trust risk\*** ? the worst case is under-trusting a correct parse. Float arithmetic is fine here (tolerance `0.005` ? accumulation error) and consistent with the savings parsers. CI is green (480 tests, ruff/mypy) on Ubuntu 3.11/3.12/3.13 + Windows.

Four things to consider ? none are blockers, but the first two are worth doing:

### 1. The "Cr" (credit-balance) case is documented but not actually handled

The regex comment lists `Total Payment Due 12,345.67 Cr` as a supported form, but the trailing `Cr` is dropped ? I verified `\_TOTAL\_DUE\_RE` captures `12,345.67` (positive) and ignores the `Cr`. So when the customer has a **\*credit balance\*** (overpaid; true due is negative), `total\_due` is stored as `+12345.67` while the actual net is ? `?12345.67` ? the cross-check can **\*never\*** match ? the statement silently stays `UNVERIFIED`. Safe direction, but the comment implies it's handled. Either negate `total\_due` when a trailing `Cr`/`CR` is present, or drop that example from the comment so it doesn't over-promise.

### 2. `\_TOTAL\_DUE\_RE` has no direct test

All five tests monkeypatch `parse\_axis\_cc` to return a fixed 4-tuple, so they exercise the cross-check arithmetic in `parse\_and\_insert` ? but the **\*regex itself is never asserted\***, and it's the fragile part (three formats, the `Cr` case, number-on-next-line, "Minimum Amount Due" vs "Total Amount Due" disambiguation). A handful of direct `\_TOTAL\_DUE\_RE` cases would lock in the documented formats ? and would have surfaced #1.

### 3. Worth a one-line note: the check assumes no carried-over balance

`?(Dr) ??(Cr)` equals "Total Amount Due" only when there's **\*no uncaptured previous/opening balance\***. A revolving balance (carried from last cycle) makes the net legitimately differ from the printed total, so a correctly-parsed statement stays `UNVERIFIED`. That's fine by design, but a comment naming the limitation would stop a future reader from "fixing" it as a bug.

### 4. Minor: hoist the `BALANCE\_TOLERANCE` import

`from .validation import BALANCE\_TOLERANCE` is imported inside `parse\_and\_insert`, but `validation` is already imported at module top (line 31: `from .validation import Confidence, parse\_amount`). No circular-import reason ? just add `BALANCE\_TOLERANCE` to that existing line for consistency.

Safe to merge as-is; I'd fold in #1 + #2 either here or as a quick follow-up.

```
===== avidullu @ 2026-06-25T17:38:51Z (id 4802475958) =====  
All four review notes addressed in e5ace4e:
```

- \*Cr case handled\*** ? `\_TOTAL\_DUE\_RE` now captures the optional `Cr` suffix (group 2); `parse\_axis\_cc` negates `total\_due` when it's present. Credit-balance statements can now graduate to RECONCILED.
- \*Direct regex tests\*** ? 5 new tests for `\_TOTAL\_DUE\_RE`: plain number, (Rs.) variant, Payment Due variant, Cr suffix capture, and negative (no-match on Minimum/other lines). These would have caught #1.
- \*Carried-over balance comment\*** ? added inline note in `parse\_and\_insert` documenting that `?(Dr)??(Cr)` = total only without a revolving balance; legitimate mismatch stays UNVERIFIED by design.
- \*Import hoisted\*** ? `BALANCE\_TOLERANCE` moved to module-level import alongside `Confidence`/`parse\_amount`.

485 tests pass, ruff/mypy clean. Ready for re-review.

```
=== #83 reviews + inline ===
=== #83 current head vs my reviewed head (75d220ca) ===
{"head":"75d220ca73580bc9b3ade1cc254986960ef7de43","state":"OPEN"}
```

#### 150. assistant (2026-06-25T17:39:48.847Z)

The author says #83's four notes were addressed in commit `e5ace4e` ? but the PR head is still `75d220ca` (what I reviewed), **not** `e5ace4e`. Same shape as #82, except here the head didn't move. Let me check whether `e5ace4e` was actually pushed to the PR branch.

#### 151. user (2026-06-25T17:39:55.459Z)

```
=== git's view of the PR branch head ===
e5ace4eabca904f1bc6eed690be5ced88b10245

=== does commit e5ace4e exist anywhere we can see? ===
commit

=== branch tip log ===
e5ace4e fix: address PR #83 review ? Cr handling, regex tests, balance note, import hoist
75d220c feat: add total-due cross-check to Axis CC parser
ec12b59 fix: address PR #80 post-merge review comments (#81)

=== which remote branches (if any) contain e5ace4e? ===
origin/feat/axis-cc-total-due

=== search ALL remote refs for the fix commit message ===
e5ace4e fix: address PR #83 review ? Cr handling, regex tests, balance note, import hoist
75d220c feat: add total-due cross-check to Axis CC parser
3ab48ff fix(parsers): Axis CC floor -- explicit "unverified" verdict, no silent 0.0 amounts
(#59)
```

#### 152. assistant (2026-06-25T17:40:14.302Z)

`gh pr view` was stale ? git (post-fetch) shows the PR head **is** `e5ace4e` ("fix: address PR #83 review ? Cr handling, regex tests, balance note, import hoist"). So the fix was pushed. This one changes parsing *logic* (Cr negation), so it needs a real correctness re-review, not just presence checks. Let me get the fix-commit diff.

#### 153. user (2026-06-25T17:40:18.015Z)

```
=== fix commit diff: 75d220c -> e5ace4e ===
diff --git a/src/muneem/parsers/axis.py b/src/muneem/parsers/axis.py
index 11455b5..ba6aa76 100644
--- a/src/muneem/parsers/axis.py
+++ b/src/muneem/parsers/axis.py
@@ -28,16 +28,18 @@ from .base import StatementParser
 from .dates import normalize_statement_date
 from .profiles import AXIS_PROFILE
 from .savings import SavingsParse, parse_savings, parse_statement_summary, txn_line_predicate
-from .validation import Confidence, parse_amount
+from .validation import BALANCE_TOLERANCE, Confidence, parse_amount

_CC_DATE_RE = re.compile(r"\d{2}/\d{2}/\d{4}") # Axis credit card: dd/mm/yyyy
_CARD_RES = [re.compile(r"(\d{4}XXX\d{4})"), re.compile(r"(\d{6}\*\d{4})")]
# Axis CC statements print the total amount due in forms like:
#   Total Amount Due           12,345.67
#   Total Amount Due (Rs.)     12,345.67
-#   Total Payment Due          12,345.67 Cr
+#   Total Payment Due          12,345.67
+# Cr/Debit suffixes (if any) are captured separately for sign handling.
_TOTAL_DUE_RE = re.compile(
```

```
- r"Total\s+(?:Amount|Payment)\s+Due\s*(?:\s*\s*)?\s*([\d,]+\.\s*\d{0,2})",
+ r"Total\s+(?:Amount|Payment)\s+Due\s*(?:\s*\s*)?\s*"
+ r"([\d,]+\.\s*\d{0,2})\s*(Cr)?",
  re.IGNORECASE,
)
```

```
@@ -189,7 +191,14 @@ def parse_axis_cc(
    card_number = _first_match(head, _CARD_RES) or "UNKNOWN"
    m = _TOTAL_DUE_RE.search(head)
    if m:
-        total_due = parse_amount(m.group(1).replace(",", ""))
+        raw = m.group(1).replace(",", "")
+        total_due = parse_amount(raw)
+        # A trailing "Cr" means the customer has a credit balance
+        # (overpaid); the real net is negative. Without the negation
+        # the cross-check would never match ? safe, but the parse
+        # would silently stay UNVERIFIED.
+        if total_due is not None and m.group(2):
+            total_due = -total_due
    for page in pdf.pages:
        page_rows, page_bad = _rows_from_cc_tables(page.extract_tables())
        rows.extend(page_rows)
@@ -235,11 +244,13 @@ class AxisCreditCardParser(StatementParser):
    return 0
```

```

# Cross-check against the printed total-due when available.
+ # Note: ?(Dr) ? ?(Cr) equals "Total Amount Due" only when there is
+ # no uncaptured carried-over/revolving balance from a prior cycle.
+ # When that assumption doesn't hold (legitimate mismatch), the parse
+ # stays UNVERIFIED ? correct by design, not a bug.
confidence = Confidence.UNVERIFIED
if total_due is not None:
    net = sum(r.amount if r.type == "Dr" else -r.amount for r in rows)
- from .validation import BALANCE_TOLERANCE
-
    if abs(net - total_due) <= BALANCE_TOLERANCE:
        confidence = Confidence.RECONCILED
```

=== CI on e5ace4e ===

```
test (ubuntu-latest, 3.11)      pass    37s
https://github.com/avidullu/muneem/actions/runs/28187232776/job/83492203684
test (ubuntu-latest, 3.12)      pass    48s
https://github.com/avidullu/muneem/actions/runs/28187232776/job/83492203719
test (ubuntu-latest, 3.13)      pass    38s
https://github.com/avidullu/muneem/actions/runs/28187232776/job/83492203739
test (windows-latest, 3.13)     pass    1m37s
https://github.com/avidullu/muneem/actions/runs/28187232776/job/83492203761
test (ubuntu-latest, 3.11)      pass    37s
https://github.com/avidullu/muneem/actions/runs/28187222630/job/83492168912
test (ubuntu-latest, 3.12)      pass    47s
https://github.com/avidullu/muneem/actions/runs/28187222630/job/83492168899
```

#### 154. assistant (2026-06-25T17:40:53.228Z)

The code fix looks correct on inspection ? but I need three things before signing off: the **\*\*regex test additions\*\*** (note #2), an empirical check of the **\*\*new regex\*\*** (esp. the ``(Cr)?`` group under IGNORECASE), and confirmation that **\*\*CI actually ran green on `e5ace4e`\*\*** (gh's view was lagging). Let me get all three.

#### 155. user (2026-06-25T17:40:55.346Z)

```
=== test diff (regex tests for note #2) ===
diff --git a/tests/parsers/test_axis.py b/tests/parsers/test_axis.py
index 77bfb82..3ed22ad 100644
```

```

--- a/tests/parsers/test_axis.py
+++ b/tests/parsers/test_axis.py
@@ -200,3 +200,31 @@ def test_cc_total_due_cross_check_stays_unverified_on_mismatch(tmp_path,
monkeyp
)
    assert n == 1 and len(stored) == 1
    assert status == "unverified" # rows stored but not trusted
+
+
+## --- _TOTAL_DUE_RE direct tests -----
+
+
+def test_total_due_re_plain_number():
+    m = axis._TOTAL_DUE_RE.search("Total Amount Due          12,345.67")
+    assert m and m.group(1) == "12,345.67" and m.group(2) is None
+
+
+def test_total_due_re_with_rs_parens():
+    m = axis._TOTAL_DUE_RE.search("Total Amount Due (Rs.)      12,345.67")
+    assert m and m.group(1) == "12,345.67"
+
+
+def test_total_due_re_payment_due_variant():
+    m = axis._TOTAL_DUE_RE.search("Total Payment Due          1,234.56")
+    assert m and m.group(1) == "1,234.56"
+
+
+def test_total_due_re_with_cr_suffix():
+    m = axis._TOTAL_DUE_RE.search("Total Payment Due          1,234.56 Cr")
+    assert m and m.group(1) == "1,234.56" and m.group(2) == "Cr"
+
+
+def test_total_due_re_no_match_on_other_lines():
+    assert axis._TOTAL_DUE_RE.search("Minimum Amount Due      500.00") is None
+    assert axis._TOTAL_DUE_RE.search("Credit Limit           100,000") is None

```

#### 156. user (2026-06-25T17:41:02.594Z)

```

=== empirical: new regex behavior ===
'Total Amount Due          12,345.67'          g1='12,345.67' g2=None    -> 12345.67
'Total Amount Due (Rs.)    12,345.67'          g1='12,345.67' g2=None    -> 12345.67
'Total Payment Due         12,345.67 Cr'        g1='12,345.67' g2='Cr'     -> -12345.67
'Total Amount Due          12,345.67 Credit'    g1='12,345.67' g2='Cr'     -> -12345.67
'Minimum Amount Due 500.00'                     NO MATCH
'Total Amount Due\n8,000.00\nMinimum Amount D' g1='8,000.00' g2=None    -> 8000.0

```

=== CI status for commit e5ace4e specifically (check-runs) ===

#### 157. assistant (2026-06-25T17:42:08.653Z)

Found something important. Let me nail down two things: a likely **regression** in the new `(Cr)?` group\*\*, and whether CI has actually reported on `e5ace4e` (gh's view is lagging).

#### 158. user (2026-06-25T17:42:19.828Z)

```

=== does the (Cr)? overmatch real CC fields? + does \b fix it? ===
C:\Users\avidu\AppData\Local\Temp\re3.py:2: SyntaxWarning: invalid escape sequence '\s'
    cur = re.compile(r"...s*([\d,]+\.\d{0,2})\s*(Cr)?".replace("...", "Total\s+(?:Amount|Payment)\s+Due"), re.IGNORECASE)
CURRENT      g2='Cr'    total_due=-12345.67    <- should be +12345.67
WITH ■       g2=None    total_due=12345.67    <- should be +12345.67

```

=== CI: workflow runs whose head is e5ace4e ===

--- (if empty above, no run has e5ace4e as head yet) ---

check-runs total\_count for e5ace4e: 0

### 159. assistant (2026-06-25T17:42:58.764Z)

Both confirmed ? and this is a real catch. Let me write up the re-review.

The new `(Cr)?` group is a **regression**: under `IGNORECASE` with no word boundary, it matches the "Cr" in **"Credit"**. On a realistic summary-box layout where "Credit Limit" follows the total-due in the extracted text (`Total Amount Due\n12,345.67\nCredit Limit\n2,00,000`), it captures `g2='Cr'` and **negates** a normal positive balance to ?12,345.67 ? cross-check fails ? stays `UNVERIFIED`. The original commit handled this case correctly. Adding `b` (`Cr\b`) fixes it (verified). Also: **CI** has not run on `e5ace4e` (check-runs count = 0; the green checks gh shows are from the prior commit).

### 160. assistant (2026-06-25T17:43:13.232Z)

## Re-review of `e5ace4e` ? 3 of 4 addressed cleanly; the Cr fix introduced a regression ??

Thanks for the quick turnaround. Three are solid:

- **#2 direct regex tests** ? ? the 5 `\_TOTAL\_DUE\_RE` cases (plain / `(Rs.)` / Payment Due / `Cr` suffix / no-match on Minimum & Credit Limit) are exactly right.
- **#3 carried-over-balance note** ? ? clear inline comment, no behavior change.
- **#4 import hoist** ? ? `BALANCE\_TOLERANCE` now alongside `Confidence`/`parse\_amount`, function-local import removed, no duplicate.

### ?? #1 (Cr handling) is logically right for a real `Cr`, but the regex over-matches "Credit"

`(Cr)?` under `re.IGNORECASE` with no trailing word boundary matches the **"Cr"** in **"Credit"**. Credit-card summary boxes almost always print **"Credit Limit"** (and "Credit Card") right next to "Total Amount Due", and `s` spans newlines, so on a normal **positive-balance** statement the regex grabs the wrong suffix and **negates** the total:

```
python
text = "Total Amount Due\n12,345.67\nCredit Limit\n2,00,000.00"
m = _TOTAL_DUE_RE.search(text)
m.group(1), m.group(2) # -> ('12,345.67', 'Cr') ? from "Credit", not a real Cr
total_due becomes -12,345.67 -> |net - total_due| ? 24,691 -> stays UNVERIFIED
```

So a correctly-parsed normal statement silently fails to reconcile ? the exact case this PR is meant to graduate to `RECONCILED`. It's still *safe* (never falsely reconciles), but it's a **regression** vs `75d220c`, which captured `12,345.67` positive and reconciled fine. The new tests miss it because they feed clean strings where the number is followed by end-of-string or a non-`cr` word.

**One-char fix** ? anchor the suffix with a word boundary:

```
python
r"...([\d,]+\.\d{0,2})s*(Cr\b)?"
```

Verified: `Cr\b` still captures a genuine `Cr` but no longer matches `Credit` (g2 ? `None` ? total stays `+12,345.67`). Worth adding the adjacency case as a test:

```
python
def test_total_due_re_does_not_match_cr_inside_credit_limit():
    m = axis._TOTAL_DUE_RE.search("Total Amount Due\n12,345.67\nCredit Limit 2,00,000")
    assert m and m.group(1) == "12,345.67" and m.group(2) is None
```

(Worth a quick sanity check against a real Axis CC statement's extracted text to confirm the adjacency ? but the over-match itself is definite, and the `b` fix is zero-risk regardless.)

### CI hasn't run on this commit yet

The checks currently shown on the PR are from the prior commit `75d220c` ? `check-runs` for `e5ace4e` is empty (your "485 pass" is the local run). I'll confirm green on `e5ace4e` once the workflow reports.

Holding sign-off on the Cr regression + pending CI. Everything else is good to go.

<!-- claude-review-bot -->

### 161. user (2026-06-25T17:43:13.258Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\pr83-rereview.md (file state is current in your context ? no need to Read it back)

### 162. user (2026-06-25T17:43:31.620Z)

<https://github.com/avidullu/muneem/pull/83#issuecomment-4802521424>

```
=== update state: #83 head=e5ace4e (re-reviewed), advance last_comment_ts ===
{
  "head": "8306ba8987f23dea28caf54560f3cfa69dab2241",
  "last_comment_ts": "2026-06-25T17:43:30Z"
}
```

### 163. assistant (2026-06-25T17:43:47.094Z)

The head moved **again** while I was reviewing ? `git rev-parse` now shows `8306ba8`, not `e5ace4e`. The author is pushing rapidly. I may have just re-reviewed a stale commit. Let me check what `8306ba8` is before anything else.

### 164. user (2026-06-25T17:43:53.197Z)

```
=== branch log now ===
8306ba8 test: add credit-balance cross-check + no-Cr regex tests
e5ace4e fix: address PR #83 review ? Cr handling, regex tests, balance note, import hoist
75d220c feat: add total-due cross-check to Axis CC parser
ec12b59 fix: address PR #80 post-merge review comments (#81)
```

=== what did 8306ba8 change vs e5ace4e? ===

```
=== timing: when was my re-review vs the new commit? ===
my re-review comment (4802521424) created_at:
2026-06-25T17:43:30Z
8306ba8 commit time:
2026-06-25T23:12:35+05:30
```

### 165. assistant (2026-06-25T17:44:13.422Z)

`8306ba8` is "test: add credit-balance cross-check + no-Cr regex tests" ? and crucially the `axis.py` diff vs `e5ace4e` is **empty**, so the regex source is unchanged: my Cr over-match finding still stands. It landed ~1 min before my re-review (author iterating independently, not yet reacting to my comment). Let me see exactly what tests it added ? whether they catch or miss the "Credit Limit" case.

### 166. user (2026-06-25T17:44:18.893Z)

```
=== files touched by 8306ba8 ===
8306ba8 test: add credit-balance cross-check + no-Cr regex tests
  tests/parsers/test_axis.py | 22 +++++
  1 file changed, 22 insertions(+)

=== test additions in 8306ba8 ===
diff --git a/tests/parsers/test_axis.py b/tests/parsers/test_axis.py
index 3ed22ad..4fafd8d 100644
```



```

--- a/tests/parsers/test_axis.py
+++ b/tests/parsers/test_axis.py
@@ -228,3 +228,25 @@ def test_total_due_re_with_cr_suffix():
    def test_total_due_re_no_match_on_other_lines():
        assert axis._TOTAL_DUE_RE.search("Minimum Amount Due          500.00") is None
        assert axis._TOTAL_DUE_RE.search("Credit Limit              100,000") is None
+
+
+def test_total_due_re_without_cr_suffix():
+    """Normal (debit-balance) case: no trailing Cr."""
+    m = axis._TOTAL_DUE_RE.search("Total Amount Due          5,000.00")
+    assert m and m.group(1) == "5,000.00" and m.group(2) is None
+
+
+def test_cc_total_due_cross_check_credit_balance_reconciles(tmp_path, monkeypatch):
+    """When the customer has a credit balance (overpaid), the Cr suffix produces
+    a negative total_due. The cross-check should match the net."""
+    txns = [
+        axis.AxisCCTxn("01/03/2026", "Payment", "Payment", 500.0, "Cr"),
+        axis.AxisCCTxn("02/03/2026", "Refund", "Refund", 100.0, "Cr"),
+    ]
+    # Net: ?500 ? 100 = ?600. Printed total = 600 Cr ? total_due = ?600
+    n, status, stored, prov = _cc_parse_and_insert(
+        tmp_path, monkeypatch, ("CARD1", txns, 0, -600.0)
+    )
+    assert n == 2 and len(stored) == 2
+    assert status == "reconciled"
+    assert prov["confidence"] == "reconciled"
+
+
+=== CI on 8306ba8 ===
+check-runs total_count: 5
+ocr: completed/skipped
+test (ubuntu-latest, 3.11): in_progress/null
+test (ubuntu-latest, 3.12): completed/success
+test (ubuntu-latest, 3.13): completed/success
+test (windows-latest, 3.13): in_progress/null

```

#### 167. assistant (2026-06-25T17:44:54.119Z)

`8306ba8` is **test-only** (axis.py byte-identical to `e5ace4e`), so my Cr over-match finding still stands ? and crucially, the two tests it adds (`test\_total\_due\_re\_without\_cr\_suffix` feeds a clean string; the credit-balance test passes a pre-negated `?600`) **don't** exercise the "Credit Limit" adjacency\*\* that triggers the bug. CI is now running on `8306ba8` (Ubuntu 3.12/3.13 green, 3.11 + Windows still going). Let me post a tight sync note so the author doesn't think the new Cr tests closed it.

#### 168. user (2026-06-25T17:45:07.597Z)

<https://github.com/avidullu/muneem/pull/83#issuecomment-4802534849>  
state[83]: {"head": "8306ba8987f23dea28caf54560f3cfa69dab2241", "last\_comment\_ts": "2026-06-25T17:45:06Z"}

#### 169. assistant (2026-06-25T17:45:26.688Z)

Reviewed #83's comments and re-reviewed the work behind them. Here's the bottom line.

#### The author responded "all four addressed in `e5ace4e`" ? I re-reviewed it

Three of the four are cleanly done: **regex tests** (#2), **carried-balance comment** (#3), and **BALANCE\_TOLERANCE` import hoist** (#4) ? all good.

But the **Cr-handling fix** (#1) introduced a regression\*\*, which I caught by running the new regex against realistic input:

- The new `(Cr)?` group under `IGNORECASE` with no word boundary matches the **Cr** in

"Credit"\*\*. CC summary boxes almost always print **"Credit Limit"** next to "Total Amount Due", and `\s*` spans newlines ? so on a *\*normal positive-balance\** statement (``Total Amount Due\n12,345.67\nCredit Limit?``) it captures ``Cr`` from "Credit" and **\*\*negates the total to ?12,345.67\*\*** ? cross-check fails ? stays ``UNVERIFIED``. That's the exact case the feature is meant to reconcile, and the original commit handled it correctly ? so it's a regression.

- It's still *\*safe\** (never falsely reconciles), but it silently defeats the feature for the common case.
- **\*\*One-char fix:\*\*** ``(Cr\b)?`` ? verified it still catches a genuine ``Cr`` but no longer matches "Credit". I gave a suggested adjacency test too.

Then the author pushed **``8306ba8``** (test-only) **\*\*mid-review** ? it adds credit-balance tests but feeds clean strings / a pre-negated total, so it **\*\*doesn't cover the over-match\*\***; the finding stands.

## Status

- Posted: [re-review](<https://github.com/avidullu/muneem/pull/83#issuecomment-4802521424>) + a [sync note](<https://github.com/avidullu/muneem/pull/83#issuecomment-4802534849>) about the racing commit.
- **\*\*Holding sign-off\*\*** until the ``Cr\b`` fix lands. CI is mid-run on ``8306ba8`` (Ubuntu 3.12/3.13 green; 3.11 + Windows in progress).
- The poll is updated (#83 head ``8306ba8``), so if the author pushes the fix, the next tick will re-review automatically ? or just tell me and I'll check immediately.

One thing worth knowing: this branch is moving fast (three pushes in a few minutes), and ``gh``'s API view lagged ``git`` twice today ? so I've been verifying head SHAs with ``git rev-parse`` after a fetch rather than trusting ``gh pr view``. If you'd rather I batch re-reviews (wait for you to say "ready") instead of chasing each push, say the word.

### 170. user (2026-06-25T17:46:17.525Z)

great catch! Can you again review the PR with a new commit. Be as thorough as you have been!

### 171. user (2026-06-25T17:46:30.019Z)

```
=== current head + recent log ===
c455994bf057392f3f936db3765d8cf3f3e70ac5
c455994 fix: _TOTAL_DUE_RE Cr suffix uses \\b word boundary to avoid 'Credit' over-match
8306ba8 test: add credit-balance cross-check + no-Cr regex tests
e5ace4e fix: address PR #83 review ? Cr handling, regex tests, balance note, import hoist
75d220c feat: add total-due cross-check to Axis CC parser
ec12b59 fix: address PR #80 post-merge review comments (#81)
```

```
=== my last-reviewed head (from state) ===
{'head': '8306ba8987f23dea28caf54560f3cfa69dab2241', 'last_comment_ts': '2026-06-25T17:45:06Z'}
```

```
=== any new author comments since 17:45:06? ===
--- avidullu @ 2026-06-25T17:45:15Z ---
Good catch ? the 'Credit' over-match is real. Fixed in c455994:
```

- Regex: `(Cr)? ? (Cr\\b)? ?` word boundary prevents matching 'Credit' in adjacent 'Credit Limit' / 'Credit Card' lines
- New test: `test_total_due_re_does_not_match_cr_inside_credit_limit` with the exact adjacency case (Total Amount Due ? number ? Credit Limit on next line)
- Existing Cr suffix test still passes (genuine trailing Cr with nothing after it)

488 tests pass, ruff/mypy clean. CI should be green shortly.

### 172. assistant (2026-06-25T17:46:57.148Z)

The author applied exactly the ``(Cr\b)?`` fix I suggested, in ``c455994``, plus the adjacency test. Let me re-review it thoroughly: confirm the diff is *\*only\** the boundary change, exercise the new regex against a broad set of cases (including ones their test doesn't cover), and check CI.

### 173. user (2026-06-25T17:47:02.410Z)

```

=== FULL diff 8306ba8 -> c455994 (source + tests) ===
diff --git a/src/muneem/parsers/axis.py b/src/muneem/parsers/axis.py
index ba6aa76..7627164 100644
--- a/src/muneem/parsers/axis.py
+++ b/src/muneem/parsers/axis.py
@@ -39,7 +39,7 @@ _CARD_RES = [re.compile(r"(\d{4}XXX\d{4})", re.compile(r"(\d{6}\*\d{4})"))]
# Cr/Debit suffixes (if any) are captured separately for sign handling.
_TOTAL_DUE_RE = re.compile(
    r"Total\s+(?:Amount|Payment)\s+Due\s*(?:\s*(Rs\.\d{2})\s*)?\s*"
-   r"([\d,]+\.\d{0,2})\s*(Cr)?",
+   r"([\d,]+\.\d{0,2})\s*(Cr\b)?",
    re.IGNORECASE,
)

diff --git a/tests/parsers/test_axis.py b/tests/parsers/test_axis.py
index 4fafd8d..6c0c295 100644
--- a/tests/parsers/test_axis.py
+++ b/tests/parsers/test_axis.py
@@ -236,6 +236,14 @@ def test_total_due_re_without_cr_suffix():
    assert m and m.group(1) == "5,000.00" and m.group(2) is None

+def test_total_due_re_does_not_match_cr_inside_credit_limit():
+    """The Cr suffix must use \b word boundary so 'Credit Limit'
+    on an adjacent line does not trigger a false credit-balance negation."""
+    text = "Total Amount Due\n12,345.67\nCredit Limit 2,00,000"
+    m = axis._TOTAL_DUE_RE.search(text)
+    assert m and m.group(1) == "12,345.67" and m.group(2) is None
+
+def test_cc_total_due_cross_check_credit_balance_reconciles(tmp_path, monkeypatch):
+    """When the customer has a credit balance (overpaid), the Cr suffix produces
+    a negative total_due. The cross-check should match the net."""

=== executable-AST: is the ONLY axis.py change the regex literal? ===
functions whose AST changed: NONE (only a module-level constant/regex literal changed)

```

#### 174. assistant (2026-06-25T17:47:12.753Z)

Clean ? the diff is *exactly* ``(Cr)?` ? `(Cr\b)?`` plus the adjacency test, and the AST check confirms *\*\*no function logic changed\*\**, only the regex literal. Now let me exercise the new regex hard, including cases beyond the author's tests (other ``cr``-words, punctuation, uppercase, genuine Cr variants), and check CI.

#### 175. user (2026-06-25T17:47:24.685Z)

```

=== comprehensive regex behavior on c455994's _TOTAL_DUE_RE ===
plain positive          g1='12,345.67' g2=None   total_due=12345.67
(Rs.) variant          g1='12,345.67' g2=None   total_due=12345.67
Payment Due variant    g1='1,234.56'  g2=None   total_due=1234.56
genuine Cr (EOL)       g1='12,345.67' g2='Cr'    total_due=-12345.67
genuine CR (uppercase) g1='12,345.67' g2='CR'    total_due=-12345.67
genuine Cr.            g1='12,345.67' g2='Cr'    total_due=-12345.67
Credit Limit adjacency g1='12,345.67' g2=None   total_due=12345.67
Credit Card adjacency g1='8,000.00'  g2=None   total_due=8000.0
Crore (false-cr word)  g1='12,345.67' g2=None   total_due=12345.67
Credited (false-cr word) g1='9,999.00' g2=None   total_due=9999.0
Dr suffix (debit)      g1='5,000.00'  g2=None   total_due=5000.0
number on next line    g1='8,000.00'  g2=None   total_due=8000.0
Minimum (must NOT match) NO MATCH

```

```

=== CI on c455994 ===
check-runs: 5 total
ocr: completed/skipped
test (ubuntu-latest, 3.11): queued/null

```

```
test (ubuntu-latest, 3.12): in_progress/null
test (ubuntu-latest, 3.13): queued/null
test (windows-latest, 3.13): completed/success
```

#### 176. assistant (2026-06-25T17:48:14.455Z)

The `(Cr\b)?` fix is **correct and robust** ? verified across every case I could construct, including ones the author's tests don't cover:

```
| Input | Result | ? |
|---|---|---|
| genuine `Cr` / `CR` / `Cr.` (EOL) | negated (credit balance) | ? |
| `?12,345.67\nCredit Limit?` | positive (no false negation) | ? |
| `?8,000.00 Credit Card?` | positive | ? |
| `Crore` / `Credited` (other cr-words) | positive | ? |
| `Dr` suffix / number-on-next-line | positive | ? |
| `Minimum Amount Due` | no match | ? |
```

AST-confirmed that **only the regex literal changed** ? no other logic touched. Windows CI is already green; Ubuntu is mid-run. Let me watch it to completion so I can give a definitive sign-off rather than a conditional one.

#### 177. user (2026-06-25T17:50:17.786Z)

Command running in background with ID: blawjublu. Output is being written to:  
C:\Users\avidu\AppData\Local\Temp\claudu\C--Users-avidu-Projects-badminton-highlight-indexer\6790052a-f068-488a-9cc6-95b0c547eeb1\tasks\blawjublu.output. You will be notified when it completes. To check interim output, use Read on that file path.

#### 178. user (2026-06-25T17:50:41.035Z)

```
1      watching ci run 28189109387 for c455994 ...
2
```

#### 179. assistant (2026-06-25T17:50:54.452Z)

Thorough review is done ? the fix is correct and comprehensively verified. I'm watching CI run `28189109387` to completion in the background so I can give you a definitive sign-off rather than a conditional one (Windows already green; Ubuntu 3.11/3.12/3.13 finishing). I'll post the LGTM to the PR and confirm here the moment it lands green ? or flag immediately in the unlikely event Ubuntu surfaces anything Windows didn't.

Summary so far on `c455994`:

- **Regression fixed correctly** ? `(Cr)?` ? `(Cr\b)?`, the exact boundary fix; genuine `Cr`/`CR`/`Cr.` still negate, but `Credit Limit`/`Credit Card`/`Crore`/`Credited` no longer false-trigger.
- **AST-confirmed surgical** ? only the regex literal changed; zero function-body changes, so the cross-check arithmetic I reviewed earlier is untouched.
- **Right test added** ? `test\_total\_due\_re\_does\_not\_match\_cr\_inside\_credit\_limit` pins the exact adjacency case.
- All four original review notes now fully resolved.

Standing by for CI.

#### 180. user (2026-06-25T17:51:24.704Z)

CI is green

#### 181. user (2026-06-25T17:51:36.292Z)

```
=== confirm head still c455994 (no newer push) ===
c455994bf057392f3f936db3765d8cf3f3e70ac5
```

```
=== CI conclusions for c455994 ===
```

```
ocr: completed/skipped
test (ubuntu-latest, 3.11): queued/null
test (ubuntu-latest, 3.12): completed/success
test (ubuntu-latest, 3.13): in_progress/null
test (windows-latest, 3.13): completed/success
```

```
=== overall: any non-success (excl. skipped)? ===
0
```

#### 182. assistant (2026-06-25T17:51:51.789Z)

The user says green, but the check-runs API still shows 3.11 `queued` and 3.13 `in\_progress` ? likely the gh-API lag I've hit all session, or a duplicate run. Let me reconcile against the actual workflow runs and the watch result before I sign off.

#### 183. user (2026-06-25T17:51:58.525Z)

```
=== all workflow runs whose head is c455994 ===
run 28189109387 [push] ci: completed/success
```

```
=== background watch (run 28189109387) final output ===
{"conclusion":"success","jobs":[{"conclusion":"success","name":"test (ubuntu-latest,
3.11)"}, {"conclusion":"success","name":"test (ubuntu-latest,
3.13)"}, {"conclusion":"success","name":"test (windows-latest,
3.13)"}, {"conclusion":"success","name":"test (ubuntu-latest,
3.12)"}, {"conclusion":"skipped","name":"ocr"}], "status":"completed"}
```

#### 184. assistant (2026-06-25T17:52:16.312Z)

Definitive: run **28189109387** (`ci`, push) is **completed/success** ? all four test jobs (Ubuntu 3.11/3.12/3.13 + Windows 3.13) passed, ocr skipped, and the watch exited 0. The check-runs API was just lagging again. You're right, it's green. Let me post the final sign-off.

#### 185. user (2026-06-25T17:52:30.582Z)

```
https://github.com/avidullu/muneem/pull/83#issuecomment-4802600866
state[83]: {"head": "c455994bf057392f3f936db3765d8cf3f3e70ac5", "last_comment_ts":
"2026-06-25T17:52:29Z"}
```

#### 186. assistant (2026-06-25T17:52:40.366Z)

```
Done ? [final LGTM posted on
#83](https://github.com/avidullu/muneem/pull/83#issuecomment-4802600866).
```

**#83** is fully resolved: the `(Cr\b)?` fix in `c455994` is correct and comprehensively verified (genuine credit balances still negate; `Credit Limit`/`Credit Card`/`Crore`/`Credited` no longer false-trigger), the change is AST-surgical (only the regex literal moved), the right adjacency test is in, and CI is green on the real run (28189109387, 488 tests). All four of my original notes are closed. Clean to merge.

That closes the loop on the catch ? the over-match would have silently kept normal positive-balance statements at `UNVERIFIED`, and it's now fixed with a one-char boundary plus a regression test.

```
**Current muneem queue:**
- #82 ? LGTM (split + security-comment restoration verified)
- #83 ? LGTM (feature + Cr fix verified) ? just now
- #84 ? LGTM (docs archive)
```

All three reviewed, green, and signed off. The comment-aware poll is still live (state tracks #83 at `c455994`), so if you push more or respond, the next tick catches it ? or just ping me. Nice work on the fast iteration.