

claude-windows / 6e9d0958-a2ec-4a8e-bb2d-6cce9833c527

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\6e9d0958-a2ec-4a8e-bb2d-6cce9833c527\subagents\workflows\wf_8d68d336-288\agent-aa4401155c0c3de0e.jsonl`  
- SHA-256: `1c4f7c87d5f18cb4d25bbf51d20cc4c5ccd859eee7916f966c3cb31ef70f4170`  
- Source modified: `2026-07-03T19:51:19+00:00`  
- Imported at: `2026-07-05T16:48:26+00:00`  
- project: `wf_8d68d336-288`  
- session_id: `6e9d0958-a2ec-4a8e-bb2d-6cce9833c527`
```

Transcript

1. user (2026-07-03T19:48:21.197Z)

Repo (my in-progress optimization branch): C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree

Key files to Read:

- backend/pipeline/detectors/wasb_infer.py ? the inference loop (DataLoader ? detector.run_tensor ? tracker). Contains the streaming SequentialFrameStore, the _PrefetchIter I just added, build_windows, and the tracker re-run loop.
- backend/pipeline/detectors/native_wasb_runner.py ? builds the wasb_infer.py argv on the GPU box (native runner). Has device/stream_video/batch_size/prefetch_batches config.
- backend/config/models.py (WasbIndexConfig) and backend/config/presets.py (cloud-serving preset).
- deploy/cloudrun/Dockerfile ? the WASB conda env is torch 1.11.0+cu113 (py3.8), a SEPARATE env from the app's torch. The detector shells out to it as a subprocess.

The WASB model itself lives at /opt/models/WASB-SBDT (nttcom/WASB-SBDT) ? you can't read it, but reason from how wasb_infer.py calls build_detector(cfg) / detector.run_tensor(imgs, trans) and build_tracker(cfg).

MEASURED BASELINE (first real L4 cloud run, 2026-07-03): a ~7.5-min 60fps badminton clip took Phase-2 (WASB detection) 727.9s (~0.6x realtime); GPU utilization sat <1% the whole time. Hardware: NVIDIA L4 (Ada sm_89, 23GB), 8 vCPU, 32GB. Streaming mode forces DataLoader num_workers=0 (single-process), so CPU batch assembly (cv2 decode + PIL + ToTensor) serializes with GPU inference.

CRITICAL CONSTRAINT ? the D5/D16 byte-parity guarantee: the repo REQUIRES that the rally/window OUTPUT (the candidate windows and trajectory CSV) stay bit-identical / deterministic across runs and across the disk-frames vs streaming paths. This is asserted by a real parity test. An optimization that changes detector NUMERICS (FP16, TF32, cudnn.benchmark algo selection, NVDEC vs CPU decode) risks changing the output and breaking that guarantee. Batch size and prefetch do NOT change per-frame numerics (same forward, just grouped/overlapped) so they are parity-safe. Already implemented on this branch: indexing.wasb.batch_size + prefetch_batches knobs, the cloud-serving preset opts into batch 64 + prefetch 4.

You are a CUDA/PyTorch performance engineer. Investigate ONE optimization lever in depth by READING the actual code. Lever "torch-arch-build": Investigate bumping the WASB conda env from torch 1.11.0+cu113 to a cu12x build with NATIVE Ada (sm_89) kernels. The current cu113 wheel has no sm_89 cubin, so the first CUDA launch triggers a PTX?SASS JIT compile (the <1% util for the first ~90s of every ephemeral container run). Determine: what is the OLDEST torch build with real sm_89 kernels that still runs on Python 3.8 (the WASB env pin)? Is torch 1.11 even required by WASB-SBDT or can it go higher? Does a version bump change conv/heatmap NUMERICS (parity) or only the compiled kernel (safe)? Quantify the expected win (JIT elimination + faster Ada kernels).

Read the named files, reason from the real code, and return the structured finding. Be concrete about code_locations (file:function). Be honest about parity_risk ? the byte-parity guarantee is load-bearing.

2. assistant (2026-07-03T19:48:24.366Z)

I'll investigate the torch-arch-build lever by reading the actual code and reasoning about the version constraints.

3. user (2026-07-03T19:48:27.242Z)

```
1      """WASB shuttle-trajectory inference on an arbitrary frame sequence (runs in WSL).
2
3      WASB-SBDT only ships *dataset* evaluation; this is a thin wrapper that reuses its
4      own building blocks (ImageDataset transform + detector + online tracker) to run on
5      any folder of extracted frames and emit a trajectory CSV in the exact shape the
6      indexer's `tracknet_runner.parse_trajectory_csv` already consumes:
7
8          Frame,Visibility,X,Y
9
10     It must run INSIDE the WSL `wasb` conda env, from the WASB-SBDT `src/` dir (it
11     imports WASB's `detectors`/`trackers`/`dataloaders`). The Windows-side adapter
12     (`wasb_runner.py`) stages this file + the frames into WSL and invokes it.
13
14     Resumable across reboots
15     -----
16     The GPU detector pass (one forward per sliding window ? ~21k for a 6-min clip) is
17     the slow, expensive part; the CPU tracker pass that turns detections into a
18     trajectory is cheap and *stateful* (it must see every frame in order, so it can't
19     resume mid-stream). We exploit that split:
20
21     * The detector pass writes its per-window raw outputs incrementally to
22       ``<cache>/det_raw.jsonl`` (flushed+fsync'd every ``--flush-every`` windows) and
23       records progress in ``<cache>/manifest.json`` (atomic write). A reboot loses at
24       most ``--flush-every`` windows of GPU work instead of the whole pass.
25     * On restart we replay the cached windows, skip the DataLoader ahead to the first
26       un-cached window, and continue. Once every window is cached, the detector pass
27       is skipped entirely and we just re-run the (cheap, deterministic) tracker over
28       the full ordered cache -> trajectory CSV. So a lost trajectory CSV costs seconds,
29       not the GPU hour.
30
31     The cache is keyed by the output path (``<out>_wasbcache/`` by default), lives next
32     to the output (NOT in %TEMP%), and is invalidated automatically if the frames dir,
33     weights, sport, window size, or frame count change.
34
35     Usage (in WSL, from ~/models/WASB-SBDT/src):
36         python wasb_infer.py --frames_dir /path/frames --weights
37         ../pretrained_weights/wasb_badminton_best.pth.tar \
38         --out /path/traj.csv [--sport badminton] [--limit N] [--cache-dir DIR] [--flush-
39         every 1000] [--fresh]
40         """
41
42     import argparse
43     import csv
44     import glob
45     import json
46     import logging
47     import os
48     import os.path as osp
49     from collections import defaultdict
50     from contextlib import contextmanager
51
52     logger = logging.getLogger(__name__)
```

```
52     CACHE_VERSION = 1
```

```

53     _DET_FILE = "det_raw.jsonl"
54     _MANIFEST_FILE = "manifest.json"
55
56
57     def _build_cfg(weights: str, sport: str, device: str = "cuda"):
58         """Compose the WASB eval config via Hydra, overriding model/weights/device.
59
60         ``device`` is ``cuda`` (default ? the historical WSL behaviour, byte-parity
61         preserved),
62         ``mps`` (Apple) or ``cpu``. Only the single-GPU CUDA path requests
63         ``runner.gpus=[0]``;
64         cpu/mps run with no GPU list so the detector places tensors on ``device``.
65         """
66         from hydra import compose, initialize
67
68         gpus = "[0]" if device == "cuda" else "[]" # single GPU on cuda; none for cpu/mps
69         with initialize(config_path="configs", version_base=None):
70             cfg = compose(
71                 config_name="eval",
72                 overrides=[
73                     f"dataset={sport}",
74                     "model=wasb",
75                     f"detector.model_path={weights}",
76                     f"runner.device={device}",
77                     f"runner.gpus={gpus}", # default config assumes multiple GPUs; pin to
78                     this box
79                 ],
80             )
81             return cfg
82
83     def _frame_id(path: str) -> int:
84         """Best-effort integer frame id from a frame filename (e.g. frame_000180.png ->
85         180)."""
86         stem = osp.splitext(osp.basename(path))[0]
87         digits = "".join(ch for ch in stem if ch.isdigit())
88         return int(digits) if digits else -1
89
90     # ----- #
91     # Resumable detector cache (pure-Python, no torch/numpy ? unit-testable)
92     # ----- #
93     def _cache_paths(cache_dir: str):
94         return osp.join(cache_dir, _DET_FILE), osp.join(cache_dir, _MANIFEST_FILE)
95
96     def default_cache_dir(out_csv: str) -> str:
97         """Stable cache dir next to the trajectory output (survives reboots; not %TEMP%)."""
98         d = osp.dirname(osp.abspath(out_csv))
99         stem = osp.splitext(osp.basename(out_csv))[0]
100         return osp.join(d, stem + "_wasbcache")
101
102     def _atomic_write_text(path: str, text: str) -> None:
103         """Write then rename so a crash mid-write can't leave a half-written file."""
104         tmp = path + ".tmp"
105         with open(tmp, "w") as f:
106             f.write(text)
107             f.flush()
108             os.fsync(f.fileno())
109         os.replace(tmp, path)
110
111     def _det_plain(det) -> list:
112         """A detector candidate dict {'xy': np.array([x,y]), 'score', 'scale'} -> JSON-able

```

```

list."""
114     xy = det["xy"]
115     return [float(xy[0]), float(xy[1]), float(det["score"]), int(det["scale"])]
116
117
118     def _dets_to_tracker(plain_list):
119         """Inverse of _det_plain: rebuild the dicts the tracker expects (xy MUST be a numpy
array,
120         the tracker does vector math / np.linalg.norm on it)."""
121         import numpy as np
122
123         return [
124             {"xy": np.array([p[0], p[1]], dtype=float), "score": p[2], "scale": p[3]}
125             for p in plain_list
126         ]
127
128
129     def write_manifest(cache_dir: str, manifest: dict) -> None:
130         _, mpath = _cache_paths(cache_dir)
131         _atomic_write_text(mpath, json.dumps(manifest, indent=2))
132
133
134     def read_manifest(cache_dir: str):
135         _, mpath = _cache_paths(cache_dir)
136         if not osp.exists(mpath):
137             return None
138         try:
139             with open(mpath) as f:
140                 return json.load(f)
141         except (json.JSONDecodeError, OSError):
142             return None
143
144
145     def manifest_compatible(
146         manifest: dict,
147         *,
148         frames_dir: str,
149         weights: str,
150         sport: str,
151         frames_in: int,
152         frames_total: int,
153         device: str = "cuda",
154     ) -> bool:
155         """A cache is reusable only if the run that produced it matches this run's inputs.
156
157         ``device`` defaults to ``cuda`` so caches written before device-keying (no
158         ``device``
159         field) stay compatible with a CUDA run ? but a cpu/mps run will NOT reuse a cuda
160         cache
161         (detections can differ numerically across devices), and vice-versa.
162
163         ``frames_dir`` is the **already-canonical cache key** the caller stored in the
164         manifest
165         (``run()`` passes its ``cache_key`` here, and writes the same value into the
166         manifest):
167         ``osp.abspath(frames_dir)`` for the frames-on-disk path, or the synthetic
168         ``video:<abspath>``
169         key for the streaming path. We therefore compare it VERBATIM ? do NOT re-``abspath``
170         it here:
171         ``osp.abspath("video:/x.mp4")`` is not recognised as absolute and gets the CWD
172         prepended,
173         so a streaming cache would never match itself and every re-run would recompute
174         (issue #348).
175         """
176         if not manifest or manifest.get("version") != CACHE_VERSION:

```

```

169         return False
170     return (
171         manifest.get("frames_dir") == frames_dir
172         and manifest.get("weights") == weights
173         and manifest.get("sport") == sport
174         and manifest.get("frames_in") == frames_in
175         and manifest.get("frames_total") == frames_total
176         and manifest.get("device", "cuda") == device
177     )
178
179
180 def reset_cache(cache_dir: str) -> None:
181     """Drop any existing cache so the next run starts clean."""
182     det_jsonl, mpath = _cache_paths(cache_dir)
183     for p in (det_jsonl, mpath):
184         if osp.exists(p):
185             os.remove(p)
186
187
188 def load_and_clean_cache(cache_dir: str):
189     """Read the longest *contiguous* (w == line index) valid prefix of det_raw.jsonl,
190     rebuild the per-frame detection accumulation, and rewrite the file to exactly that
191     prefix so a trailing half-written line from a crash can't corrupt future appends.
192
193     Returns (windows_done, det_by_fid) where det_by_fid maps frame_id -> list of
194     [x, y, score, scale] candidates (accumulated across every window the frame is in,
195     in window order ? identical to a non-resumed run).
196     """
197     det_jsonl, _ = _cache_paths(cache_dir)
198     det_by_fid = defaultdict(list)
199     valid: list = []
200     if osp.exists(det_jsonl):
201         with open(det_jsonl, "r") as f:
202             for line in f:
203                 s = line.strip()
204                 if not s:
205                     continue
206                 try:
207                     obj = json.loads(s)
208                 except json.JSONDecodeError:
209                     break # truncated trailing line (crash mid-flush)
210                 if obj.get("w") != len(valid): # non-contiguous -> stop at the gap
211                     break
212                 valid.append(s)
213                 for fid, plain in obj["f"]:
214                     det_by_fid[int(fid)].extend([list(p) for p in plain])
215                 # Guarantee a clean append boundary by rewriting only the good prefix.
216                 _atomic_write_text(det_jsonl, ("\n".join(valid) + "\n") if valid else "")
217     return len(valid), det_by_fid
218
219
220 def _append_windows(fh, objs) -> None:
221     """Append window records as JSONL and durably flush (bounded loss on crash)."""
222     for o in objs:
223         fh.write(json.dumps(o, separators=(",", ":")) + "\n")
224     fh.flush()
225     os.fsync(fh.fileno())
226
227
228 def _atomic_write_trajectory(out_csv: str, rows) -> None:
229     os.makedirs(osp.dirname(osp.abspath(out_csv)), exist_ok=True)
230     tmp = out_csv + ".tmp"
231     with open(tmp, "w", newline="") as f:
232         w = csv.writer(f)
233         w.writerow(["Frame", "Visibility", "X", "Y"])

```

```

234         w.writerows(rows)
235         f.flush()
236         os.fsync(f.fileno())
237     os.replace(tmp, out_csv)
238
239
240     # ----- #
241     # Frame addressing + windowing (pure; no torch/cv2 ? unit-testable)
242     # ----- #
243     _STREAM_FRAME_FMT = "{:08d}.png"
244
245
246     def synthetic_frame_path(index: int) -> str:
247         """Stable, index-encoding frame name used by the streaming path.
248
249         It mirrors the on-disk names ``extract_frames`` writes (``{i:08d}.png``) so the
250         downstream integer-id parser (``_frame_id``) yields the SAME frame id whether the
251         frames came from disk or from the in-memory stream. The streaming image loader
252         inverts this name back to an index to fetch the decoded frame.
253         """
254         return _STREAM_FRAME_FMT.format(int(index))
255
256
257     def build_windows(frames: list, frames_in: int):
258         """Sliding windows of ``frames_in`` consecutive frame *paths* (the unit of GPU work
259         + checkpoint). Pure list math, identical for the disk and streaming paths.
260
261         Returns a list of windows, each a list of ``frames_in`` consecutive paths
262         (``[frames[i:i+frames_in] for i in range(len(frames) - frames_in + 1)]``). The
263         caller wraps each window into the ``{"frames", "annos"}`` sample shape WASB's
264         ``ImageDataset`` expects; keeping that out of here stays torch-free for unit tests.
265         """
266         return [frames[i : i + frames_in] for i in range(len(frames) - frames_in + 1)]
267
268
269     class SequentialFrameStore:
270         """In-memory sliding buffer over a forward-only frame reader, sized for the
271         detector's sliding-window access pattern (peak O(window) frames, not O(video)).
272
273         The detector DataLoader runs with ``shuffle=False`` and (in streaming mode)
274         ``num_workers=0``. ``ImageDataset.__getitem__(i)`` reads the frames of window ``i``
275         in order ? ``i, i+1, ..., i+window-1`` ? and samples are requested ``i = start,
276         start+1, ...``. So the global read sequence dips back by ``window-1`` at each window
277         boundary (``0,1,2, 1,2,3, 2,3,4, ...`` for ``window=3``); the highest index seen so
278         far never decreases. We keep only frames at or above ``max_seen - (window-1)`` (the
279         earliest index any not-yet-finished window can still ask for) and decode each source
280         frame exactly once via the forward-only ``reader(index)``.
281         """
282
283         def __init__(self, reader, total: int, window: int = 1, start_index: int = 0):
284             self._reader = reader
285             self._total = int(total)
286             self._window = max(1, int(window))
287             self._buf: dict = {}
288             # On a resumed run the detector's first needed frame is `start_index`; begin
289             # pulling there so the reader can skip (seek past) the already-cached prefix
290             # instead of decoding 0..start_index-1 just to throw them away.
291             self._next = int(start_index) # next index not yet pulled from the reader
292             self._max_seen = int(start_index) - 1
293             self._floor = int(start_index) # lowest index we still keep (never evict below)
294
295         def get(self, index: int):
296             index = int(index)
297             if index < self._floor:
298                 raise KeyError(

```

```

299         f"frame {index} already evicted (below floor {self._floor}); the access
300         "
301         f"pattern must stay within a window of the highest frame seen"
302     )
303     if index >= self._total:
304         raise KeyError(f"frame {index} out of range (total {self._total})")
305     # Pull forward from the reader until the requested index is buffered.
306     while self._next <= index:
307         self._buf[self._next] = self._reader(self._next)
308         self._next += 1
309     if index > self._max_seen:
310         self._max_seen = index
311     # Evict frames older than the earliest a still-running window could re-request:
312     # once we've seen index `m`, no future window starts before `m - (window-1)`.
313     new_floor = max(self._floor, self._max_seen - (self._window - 1))
314     for k in range(self._floor, new_floor):
315         self._buf.pop(k, None)
316     self._floor = new_floor
317     return self._buf[index]
318
319 def _index_from_synthetic_path(path: str) -> int:
320     """Invert ``synthetic_frame_path``: a frame path -> its integer source index.
321
322     Reuses ``_frame_id`` (digits-only parse) so the index and the cached frame-id stay
323     in lockstep with the on-disk naming scheme.
324     """
325     return _frame_id(path)
326
327
328 def _bgr_to_pil_rgb(frame_bgr):
329     """Convert an in-memory BGR uint8 frame (as ``cv2.VideoCapture`` yields) to the
330     EXACT
331     PIL RGB image WASB's disk path produces.
332
333     The disk path is ``frame_bgr -> cv2.imwrite(PNG) ->
334     Image.open(PNG).convert('RGB')``;
335     because PNG is lossless that round-trip equals ``cv2.cvtColor(frame_bgr,
336     COLOR_BGR2RGB)``
337     bit-for-bit (verified empirically, max|diff|=0). Returning that here makes the
338     streamed
339     frame indistinguishable from the on-disk one to everything downstream.
340     """
341     import cv2
342     from PIL import Image
343
344     return Image.fromarray(cv2.cvtColor(frame_bgr, cv2.COLOR_BGR2RGB))
345
346
347 def _make_streamed_read_image(store, real_read_image):
348     """Build the ``read_image`` replacement used during a streaming detector pass.
349
350     A synthetic frame path -> the in-memory decoded frame for its index (as a PIL RGB
351     image,
352     matching ``read_image``'s real return type); any path that doesn't parse to an index
353     falls
354     through to ``real_read_image``. Pure ? imports no WASB module ? so it is unit-
355     testable
356     without the WSL-only ``dataloaders`` package.
357     """
358
359     def _read_image(path, *args, **kwargs):
360         idx = _index_from_synthetic_path(path)
361         if idx < 0:
362             return real_read_image(path, *args, **kwargs)

```

```

356         return _bgr_to_pil_rgb(store.get(idx))
357
358     return _read_image
359
360
361     @contextmanager
362     def _streaming_read_image_patch(
363         frame_loader, total: int, window: int = 1, start_index: int = 0
364     ):
365         """Scope a shim over WASB's ``read_image`` so ``ImageDataset`` is fed in-memory
366         decoded
367         frames during the detector pass, byte-for-byte as it would over real PNGs.
368
369         WASB's ``ImageDataset.__getitem__`` reads each frame via ``read_image(path)`` ? NOT
370         ``cv2.imread``. ``read_image`` (``utils.utils``) does
371         ``Image.open(path).convert('RGB')``,
372         returning a PIL **RGB** image, after an ``osp.exists(path)`` guard. We therefore
373         feed it a
374         PIL RGB image built from the streamed BGR frame (see ``_bgr_to_pil_rgb``), which is
375         byte-identical to the disk round-trip, and the shim never touches the filesystem so
376         the
377         ``osp.exists`` guard is sidestepped for synthetic stream paths.
378
379         We patch ``dataloaders.dataset_loader.read_image`` ? the binding the consumer
380         actually
381         calls (``dataset_loader`` does ``from utils import read_image``, so the name lives
382         in that
383         module's namespace; patching ``utils.read_image`` would NOT rebind the already-
384         imported
385         reference).
386
387         ``frame_loader(index) -> BGR uint8 ndarray`` is a forward-only sequential decoder
388         wrapped
389         in a ``SequentialFrameStore`` (sliding buffer sized to ``window`` = ``frames_in``);
390         on a
391         resume we start at ``start_index`` so the decoder seeks past already-cached windows.
392         The
393         original reader is restored on exit.
394
395         """
396         from dataloaders import dataset_loader as _dl
397
398         store = SequentialFrameStore(
399             frame_loader, total, window=window, start_index=start_index
400         )
401         real_read_image = _dl.read_image
402         # Rebind read_image in the consuming module for the duration of the detector pass;
403         # restore on exit. (Plain assignment, not setattr-with-constant, to stay ruff
404         B010-clean.)
405         _dl.read_image = _make_streamed_read_image(store, real_read_image) # type:
406         ignore[assignment]
407         try:
408             yield store
409         finally:
410             _dl.read_image = real_read_image # type: ignore[assignment]
411
412
413     # ----- #
414     # Inference
415     # ----- #
416     class _PrefetchIter:
417         """Bounded producer-thread wrapper over a batch iterator (the num_workers=0
418         DataLoader).
419
420         In streaming mode the DataLoader must stay single-process (the in-memory frame store
421         +

```



```

407     the cv2.imread shim live in this process and cannot be pickled to workers), so the
408     CPU-side batch assembly (decode + PIL + ToTensor) SERIALIZES with GPU inference ?
409     measured on the first real L4 serving run (2026-07-03) this left the GPU <1%
utilized.
410     A single producer THREAD fixes the overlap without multiprocessing: cv2/PIL/numpy
411     release the GIL for the heavy parts, the store's forward-only access pattern is
412     preserved (one consumer thread iterating sequentially), and the bounded queue caps
413     memory at ``depth`` batches. A producer exception is re-raised in the consumer."""
414
415     _END = object()
416
417     def __init__(self, src, depth: int = 2):
418         import queue
419         import threading
420
421         self._q: "queue.Queue" = queue.Queue(maxsize=max(1, int(depth)))
422         self._exc: "BaseException | None" = None
423
424         def _fill() -> None:
425             try:
426                 for item in src:
427                     self._q.put(item)
428             except BaseException as e: # noqa: BLE001 - surfaced in __next__
429                 self._exc = e
430             finally:
431                 self._q.put(self._END)
432
433         self._t = threading.Thread(target=_fill, name="wasb-prefetch", daemon=True)
434         self._t.start()
435
436     def __iter__(self):
437         return self
438
439     def __next__(self):
440         item = self._q.get()
441         if item is self._END:
442             if self._exc is not None:
443                 raise self._exc
444             raise StopIteration
445         return item
446
447
448     def run(
449         frames,
450         weights: str,
451         out_csv: str,
452         sport: str = "badminton",
453         limit: int = 0,
454         batch_size: int = 8,
455         num_workers: int = 4,
456         log_every_batches: int = 50,
457         cache_dir: "str | None" = None,
458         flush_every: int = 1000,
459         fresh: bool = False,
460         frame_loader=None,
461         source_key: "str | None" = None,
462         device: str = "cuda",
463         prefetch_batches: int = 0,
464     ) -> str:
465         """Run WASB detection+tracking over an ordered frame source -> trajectory CSV.
466
467         Two equivalent ways to supply pixels (output is bit-identical between them):
468
469         * **frames-on-disk** (default): ``frames`` is a directory path string; frames are
470         globbed (``*.png``/``*.jpg``) and ``ImageDataset`` reads each file via

```

```

``cv2.imread``.
471     * **streaming** (fast path): ``frames`` is a pre-built ordered list of (synthetic)
472     frame paths and ``frame_loader(index) -> BGR uint8 ndarray`` supplies the decoded
473     pixels in-memory. We scope a shim over WASB's ``read_image`` (the actual per-frame
474     reader, which returns a PIL RGB image) over the DataLoader pass so every other
line of
475     ``ImageDataset.__getitem__`` runs UNCHANGED ? the tensor the detector sees is
identical
476     to the disk round-trip (``cv2.imread`` PNG -> ``Image.open().convert('RGB')`` ==
477     ``cvtColor(bgr, BGR2RGB)``), verified byte-identical). Streaming forces
``num_workers=0``
478     (the in-process shim + buffer can't cross worker processes).
479
480     ``source_key`` overrides the cache-keying identity (defaults to the abspath of the
481     frames dir); the streaming path passes a stable ``video:<abspath>`` key so a resume
482     after reboot still matches.
483     """
484     import time
485
486     import torch
487     from dataloaders import build_img_transforms, build_seq_transforms
488     from dataloaders.dataset_loader import ImageDataset
489     from torch.utils.data import DataLoader
490
491     # NB: this is WASB's OWN local `trackers` module (from its repo, on sys.path inside
the
492     # `wasb` conda env), NOT the Roboflow `trackers` package. The `type: ignore`
suppresses
493     # the mypy false positive where the main env resolves `trackers` to the Roboflow
package
494     # (which has no `build_tracker`). There is NO runtime collision ? this runs in a
separate
495     # subprocess/conda env. Do NOT "clean up" the ignore without removing the Roboflow
dep.
496     from trackers import build_tracker # type: ignore[attr-defined]
497     from utils import Center
498
499     from detectors import build_detector
500
501     streaming = frame_loader is not None
502
503     cfg = _build_cfg(weights, sport, device)
504     frames_in = int(cfg["model"]["frames_in"])
505     input_wh = (int(cfg["model"]["inp_width"]), int(cfg["model"]["inp_height"]))
506     output_wh = (int(cfg["model"]["out_width"]), int(cfg["model"]["out_height"]))
507
508     if streaming:
509         # `frames` is already the ordered list of (synthetic) frame paths.
510         if not isinstance(frames, (list, tuple)):
511             raise SystemExit(
512                 "streaming mode requires `frames` to be a list of frame paths"
513             )
514         frames = list(frames)
515         frames_dir = source_key or "stream"
516     else:
517         frames_dir = frames
518         frames = sorted(
519             glob.glob(osp.join(frames_dir, "*.png"))
520             + glob.glob(osp.join(frames_dir, "*.jpg"))
521         )
522     if limit:
523         frames = frames[:limit]
524     if len(frames) < frames_in:
525         where = frames_dir if not streaming else "video stream"
526         raise SystemExit(f"need >= {frames_in} frames, found {len(frames)} in {where}")

```

```

527
528     # Sliding windows of `frames_in` consecutive frames (the unit of GPU work +
checkpoint).
529     samples = []
530     for window in build_windows(frames, frames_in):
531         annos = [
532             {"frame_path": p, "center": Center(is_visible=False, x=-1.0, y=-1.0)}
533             for p in window
534         ]
535         samples.append({"frames": window, "annos": annos})
536     windows_total = len(samples)
537
538     # ---- cache setup / resume decision ----- #
539     if cache_dir is None:
540         cache_dir = default_cache_dir(out_csv)
541         os.makedirs(cache_dir, exist_ok=True)
542         det_jsonl, _ = _cache_paths(cache_dir)
543         cache_key = source_key if source_key is not None else osp.abspath(frames_dir)
544         manifest = None if fresh else read_manifest(cache_dir)
545         resume = manifest_compatible(
546             manifest or {},
547             frames_dir=cache_key,
548             weights=weights,
549             sport=sport,
550             frames_in=frames_in,
551             frames_total=len(frames),
552             device=device,
553         )
554         if resume:
555             windows_done, det_by_fid = load_and_clean_cache(cache_dir)
556             logger.info(
557                 f"resuming from cache: {windows_done}/{windows_total} windows already
detected"
558             )
559         else:
560             if manifest is not None:
561                 logger.info("cache incompatible with this run -> starting fresh")
562                 reset_cache(cache_dir)
563                 windows_done, det_by_fid = 0, defaultdict(list)
564
565         base_manifest = {
566             "version": CACHE_VERSION,
567             "frames_dir": cache_key,
568             "weights": weights,
569             "sport": sport,
570             "frames_in": frames_in,
571             "frames_total": len(frames),
572             "device": device,
573             "windows_total": windows_total,
574             "windows_done": windows_done,
575         }
576         write_manifest(cache_dir, base_manifest)
577
578     # ---- detector pass over the un-cached windows ----- #
579     remaining = samples[windows_done:]
580     if remaining:
581         _, transform_test = build_img_transforms(cfg)
582         try:
583             _, seq_transform_test = build_seq_transforms(cfg)
584         except Exception:
585             seq_transform_test = None
586         ds = ImageDataset(
587             cfg,
588             remaining,
589             input_wh,

```

```

590         output_wh,
591         transform=transform_test,
592         seq_transform=seq_transform_test,
593         is_train=False,
594     )
595     # Streaming forces single-process loading: the in-memory frame store + the
596     # cv2.imread shim live in THIS process and can't be pickled to DataLoader
workers.
597     loader_workers = 0 if streaming else num_workers
598     loader = DataLoader(
599         ds, batch_size=batch_size, shuffle=False, num_workers=loader_workers
600     )
601     detector = build_detector(cfg)
602
603     from contextlib import ExitStack
604
605     t0 = time.time()
606     done = windows_done
607     win_idx = windows_done
608     log_every = max(1, log_every_batches)
609     flush_n = max(1, flush_every)
610     pending = []
611     logger.info(
612         f"detecting on {len(remaining)} remaining windows "
613         f"(total {windows_total}, batch={batch_size}, workers={loader_workers}, "
614         f"source={'video-stream' if streaming else 'frames-dir'})..."
615     )
616     det_f = open(det_jsonl, "a")
617     try:
618         with ExitStack() as stack:
619             stack.enter_context(torch.no_grad())
620             # In streaming mode, route ImageDataset's per-frame `cv2.imread(path)`
to the
621             # in-memory decoded frame for that synthetic path; every other line of
622             # __getitem__ runs unchanged so the tensor matches the PNG round-trip
exactly.
623             # The detector's first needed frame is `windows_done` (window i starts
at
624             # frame i), so prime the store there for a resume.
625             if streaming:
626                 stack.enter_context(
627                     _streaming_read_image_patch(
628                         frame_loader,
629                         len(frames),
630                         window=frames_in,
631                         start_index=windows_done,
632                     )
633                 )
634             # Overlap CPU batch assembly with GPU inference (streaming/workers=0
only ?
635             # a multi-worker DataLoader already overlaps via its worker processes).
636             batches = (
637                 _PrefetchIter(loader, depth=prefetch_batches)
638                 if prefetch_batches and loader_workers == 0
639                 else loader
640             )
641             for bi, batch in enumerate(batches):
642                 imgs, _hms, trans = batch[0], batch[1], batch[2]
643                 img_paths = [
644                     list(t) for t in batch[-1]
645                 ] # [frame_in][batch] -> path
646                 batch_results, _ = detector.run_tensor(imgs, trans)
647                 nb = imgs.shape[0]
648                 for ib in range(nb):
649                     frames_payload = []

```

```

650         for ie in sorted(batch_results[ib].keys()):
651             p = img_paths[ie][ib]
652             fid = _frame_id(p)
653             plain = [_det_plain(d) for d in batch_results[ib][ie]]
654             frames_payload.append([fid, plain])
655             det_by_fid[fid].extend(plain)
656         pending.append({"w": win_idx, "f": frames_payload})
657         win_idx += 1
658     done += nb
659     if len(pending) >= flush_n or done >= windows_total:
660         _append_windows(det_f, pending)
661         pending = []
662         base_manifest["windows_done"] = done
663         write_manifest(cache_dir, base_manifest)
664     if (bi + 1) % log_every == 0 or done >= windows_total:
665         el = time.time() - t0
666         new_done = done - windows_done
667         rate = new_done / el if el > 0 else 0.0
668         eta = (windows_total - done) / rate if rate > 0 else 0.0
669         logger.info(
670             f"{done}/{windows_total} "
671             f"({100 * done // max(1, windows_total)}%) | {rate:.1f}
win/s | "
672             f"elapsed {el:.0f}s | ETA {eta:.0f}s"
673         )
674     if pending:
675         _append_windows(det_f, pending)
676         base_manifest["windows_done"] = done
677         write_manifest(cache_dir, base_manifest)
678     finally:
679         det_f.close()
680         logger.info(f"detection done in {time.time() - t0:.0f}s; running tracker...")
681 else:
682     logger.info("detector cache already complete -> tracker-only re-run")
683
684 # ---- tracker re-run over the full ordered cache (cheap, deterministic) ---- #
685 tracker = build_tracker(cfg)
686 tracker.refresh()
687 rows = []
688 for fid in sorted(det_by_fid.keys()):
689     r = tracker.update(_dets_to_tracker(det_by_fid[fid]))
690     rows.append((fid, 1 if r.get("visi") else 0, r.get("x", -1), r.get("y", -1)))
691
692 _atomic_write_trajectory(out_csv, rows)
693 visible = sum(1 for _, v, *_ in rows if v)
694 logger.info(f"wrote {len(rows)} rows -> {out_csv}")
695 logger.info(f"visible-shuttle frames: {visible}/{len(rows)}")
696 return out_csv
697
698
699 def extract_frames(video_path: str, frames_dir: str) -> str:
700     """Decode a video to PNG frames named by source frame index (00000000.png ...).
701
702     Idempotent: if the dir already holds the expected number of frames (matching the
703     container's reported frame count) we skip re-decoding ? re-extraction after a
704     reboot is wasteful. NOTE: cv2 PNG extraction is slow; for big clips prefer
705     extracting with ffmpeg on the host and passing --frames_dir.
706     """
707     import cv2
708
709     os.makedirs(frames_dir, exist_ok=True)
710     cap = cv2.VideoCapture(video_path)
711     if not cap.isOpened():
712         raise SystemExit(f"cannot open video: {video_path}")
713     expected = int(cap.get(cv2.CAP_PROP_FRAME_COUNT)) or 0

```

```

714     existing = len(glob.glob(osp.join(frames_dir, "*.png")))
715     if expected > 0 and existing >= expected:
716         cap.release()
717         logger.info(f"reusing {existing} already-extracted frames -> {frames_dir}")
718         return frames_dir
719     i = 0
720     while True:
721         ok, frame = cap.read()
722         if not ok:
723             break
724         cv2.imwrite(osp.join(frames_dir, f"{i:08d}.png"), frame)
725         i += 1
726     cap.release()
727     logger.info(f"extracted {i} frames -> {frames_dir}")
728     return frames_dir
729
730
731     def _probe_frame_count(video_path: str) -> int:
732         """Container-reported frame count via cv2 (0 if unknown). Used to size the
733         stream."""
734         import cv2
735         cap = cv2.VideoCapture(video_path)
736         if not cap.isOpened():
737             raise SystemExit(f"cannot open video: {video_path}")
738         n = int(cap.get(cv2.CAP_PROP_FRAME_COUNT) or 0)
739         cap.release()
740         return n
741
742
743     def make_video_frame_loader(video_path: str):
744         """Return ``(frame_loader, count_hint, release)`` for output-preserving streaming.
745
746         ``frame_loader(index) -> BGR uint8 ndarray`` decodes the video forward-only with one
747         long-lived ``cv2.VideoCapture``. It MUST be called with non-decreasing indices (the
748         detector's sliding-window order); it reads sequentially and only uses a frame-seek
749         to
750         skip forward when a *resume* starts past the current position. Each ``cap.read()``
751         returns exactly the BGR uint8 array that ``cv2.imwrite``/``cv2.imread`` of a PNG
752         would
753         yield (PNG is lossless), so the detector sees identical pixels to the disk path.
754
755         Returns the count hint from the container so the caller can build the frame list;
756         ``release`` closes the capture.
757         """
758         import cv2
759         cap = cv2.VideoCapture(video_path)
760         if not cap.isOpened():
761             raise SystemExit(f"cannot open video: {video_path}")
762         count_hint = int(cap.get(cv2.CAP_PROP_FRAME_COUNT) or 0)
763         state = {"pos": 0} # next index cap.read() will return
764
765         def frame_loader(index: int):
766             index = int(index)
767             if index < state["pos"]:
768                 raise RuntimeError(
769                     f"streaming decode is forward-only; got index {index} after
770                     {state['pos']}")
771             if index > state["pos"]:
772                 # Forward skip only happens when resuming past cached windows. Seek once.
773                 cap.set(cv2.CAP_PROP_POS_FRAMES, index)
774                 state["pos"] = index
775             ok, frame = cap.read()

```

```

775         if not ok:
776             raise RuntimeError(f"failed to decode frame {index} from {video_path}")
777         state["pos"] += 1
778         return frame
779
780     def release():
781         cap.release()
782
783     return frame_loader, count_hint, release
784
785
786 def run_video_streaming(
787     video_path: str,
788     weights: str,
789     out_csv: str,
790     sport: str = "badminton",
791     limit: int = 0,
792     batch_size: int = 8,
793     log_every_batches: int = 50,
794     cache_dir: "str | None" = None,
795     flush_every: int = 1000,
796     fresh: bool = False,
797     device: str = "cuda",
798     prefetch_batches: int = 0,
799 ) -> str:
800     """Fast path: decode ``video_path`` on the fly and feed frames straight to the
801     detector, WITHOUT extracting every frame to a PNG first (the disk round-trip that
802     dominates a long clip). Output is bit-identical to the frames-on-disk path; see
803     ``run``'s docstring for why.
804     """
805     frame_loader, count_hint, release = make_video_frame_loader(video_path)
806     try:
807         if count_hint <= 0:
808             raise SystemExit(
809                 f"could not determine frame count for {video_path}; cannot stream "
810                 f"(fall back to frame extraction with --frames_dir)"
811             )
812         frame_list = [synthetic_frame_path(i) for i in range(count_hint)]
813         source_key = "video:" + osp.abspath(video_path)
814         return run(
815             frame_list,
816             weights,
817             out_csv,
818             sport,
819             limit,
820             batch_size=batch_size,
821             num_workers=0,
822             log_every_batches=log_every_batches,
823             cache_dir=cache_dir,
824             flush_every=flush_every,
825             fresh=fresh,
826             frame_loader=frame_loader,
827             source_key=source_key,
828             device=device,
829             prefetch_batches=prefetch_batches,
830         )
831     finally:
832         release()
833
834
835 def main():
836     ap = argparse.ArgumentParser()
837     ap.add_argument("--frames_dir", help="folder of pre-extracted frames")
838     ap.add_argument(
839         "--video", help="video file; frames are extracted internally (in WSL)"
840     )

```

```

840     )
841     ap.add_argument(
842         "--frames_out_dir", help="where to extract frames when --video is used"
843     )
844     ap.add_argument(
845         "--stream-video",
846         action="store_true",
847         help="FAST PATH: decode --video on the fly and feed frames straight to "
848         "the detector (no PNG extraction). Output-equivalent to the disk path.",
849     )
850     ap.add_argument("--weights", required=True)
851     ap.add_argument("--out", required=True)
852     ap.add_argument("--sport", default="badminton")
853     ap.add_argument(
854         "--device",
855         default="cuda",
856         choices=["cuda", "mps", "cpu"],
857         help="torch device for the detector (default cuda = the WSL parity path)",
858     )
859     ap.add_argument(
860         "--limit", type=int, default=0, help="cap #frames (for quick tests)"
861     )
862     ap.add_argument("--batch-size", type=int, default=8)
863     ap.add_argument("--num-workers", type=int, default=4)
864     ap.add_argument(
865         "--prefetch-batches",
866         type=int,
867         default=0,
868         help="streaming mode: overlap CPU batch assembly with GPU inference via a "
869         "bounded producer thread of this queue depth (0 = off, the parity default)",
870     )
871     ap.add_argument(
872         "--log-every-batches", type=int, default=50, help="progress log cadence"
873     )
874     ap.add_argument(
875         "--cache-dir",
876         default=None,
877         help="resumable detector cache dir (default: <out>_wasbcache next to --out)",
878     )
879     ap.add_argument(
880         "--flush-every",
881         type=int,
882         default=1000,
883         help="fsync the detector cache every N windows (bounds reboot loss)",
884     )
885     ap.add_argument(
886         "--fresh", action="store_true", help="ignore any existing cache and start over"
887     )
888     args = ap.parse_args()
889
890     if args.stream_video:
891         if not args.video:
892             raise SystemExit("--stream-video requires --video")
893         run_video_streaming(
894             args.video,
895             args.weights,
896             args.out,
897             args.sport,
898             args.limit,
899             batch_size=args.batch_size,
900             log_every_batches=args.log_every_batches,
901             cache_dir=args.cache_dir,
902             flush_every=args.flush_every,
903             fresh=args.fresh,
904             device=args.device,

```



```

905         prefetch_batches=args.prefetch_batches,
906     )
907     return
908
909     frames_dir = args.frames_dir
910     if args.video:
911         frames_dir = extract_frames(
912             args.video, args.frames_out_dir or (args.video + "_frames")
913         )
914     if not frames_dir:
915         raise SystemExit("provide --frames_dir or --video")
916     run(
917         frames_dir,
918         args.weights,
919         args.out,
920         args.sport,
921         args.limit,
922         batch_size=args.batch_size,
923         num_workers=args.num_workers,
924         log_every_batches=args.log_every_batches,
925         cache_dir=args.cache_dir,
926         flush_every=args.flush_every,
927         fresh=args.fresh,
928         device=args.device,
929     )
930
931
932 if __name__ == "__main__":
933     main()
934

```

4. user (2026-07-03T19:48:27.876Z)

```

1     """Linux-native WASB shuttle-trajectory runner (no WSL).
2
3     The de-risk goal the ``DetectorRunner`` ABC was built for (see ``base.py`` ?
4     "Optional Linux-native path"; ``docs/DECOUPLED_COMPUTE/04-BUILD-AND-TEST-PLAN.md``
M3.5).
5     It is the sibling of :class:`~backend.pipeline.detectors.wasb_runner.WasbRunner` for a
6     Linux GPU box (GCP / a SaaS worker): it runs the same verified ``wasb_infer.py``
core
7     that the WSL adapter runs, so the trajectory CSV is identical (to a sub-pixel float
tolerance;
8     see the acceptance gate below) ? but WITHOUT the WSL transport:
9
10     * NO ``wslCommandMixin`` (no ``wsl -d <distro> bash -c``), NO ``conda activate`` (the
11     ``wasb`` env's python is invoked by path), NO ``/mnt/c`` path translation
12     (``to_wsl_mnt_path``). Paths are native; the video is already local on the box.
13     * ``device`` is ``cuda`` (default ? the WSL parity path), ``mps``, ``cpu``, or
``auto``
14     (M4: cuda-if-available-else-cpu ? the CPU fallback for a no-GPU box; explicit
``cuda`` stays
15     strict); threaded into ``wasb_infer.py --device`` (and so the Hydra
``runner.device`` override).
16     * weights / repo / python / scratch come from the environment first
17     (``WASB_WEIGHTS`` / ``WASB_REPO_DIR`` / ``WASB_PYTHON`` / ``WASB_SCRATCH``), then
18     ``indexing.wasb.*`` config, then a default ? so a box needs no config edits.
19
20     It emits the identical ``Frame,Visibility,X,Y`` CSV (same ``{stem}__{vid12}_wasb.csv``
21     name as the WSL runner) and reuses the model-agnostic ``parse_trajectory_csv`` /
22     ``trajectory_to_action_windows`` from ``tracknet_runner``, so the trajectory hybrids and
23     the windowing brain are unchanged. Selected via the existing ``detector_impl`` seam
24     (``trajectory_hybrid._resolve_runner`` ? ``_build_native_runner``);
``detector_impl="native"``.
25

```

```

26     **Acceptance gate (parity, NOT strict byte-equality):** the SAME clip through the WSL
runner
27     (Windows) and this runner (Linux, ``device=cuda``) yields the same trajectory CSV + the
same
28     candidate windows **within a small float tolerance (~1e-3 px)**. cuDNN is non-
deterministic
29     ACROSS GPUs, so a sub-pixel diff is expected on different hardware ? but it is byte-
exact
30     run-to-run on a FIXED GPU. ? Validated 2026-06-20 on an RTX 2070 (via WSL): native
streaming
31     vs the cached WSL-disk CSV matched **1194/1195** frames (the lone diff a ~5e-5 px edge
32     artifact), and two native runs were byte-identical (deterministic). This is a hardware
check
33     (a real GPU + the ``wasb`` env), so the offline, subprocess-mocked suite asserts the
contract,
34     not the numbers.
35     """
36
37     import logging
38     import os
39     import shutil
40     import subprocess
41     from dataclasses import dataclass
42     from typing import List, Optional, Tuple
43
44     from backend.config.models import IndexingConfig
45     from backend.pipeline.detectors.base import DetectorRunner
46     from backend.pipeline.detectors.device import AUTO, VALID_DEVICES, DeviceContext
47
48     # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical to
WSL/TrackNet.
49     from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
50
51     logger = logging.getLogger(__name__)
52
53     # Path to the inference core we copy into the WASB repo's src/ (so it can import WASB's
54     # detectors/trackers/dataloaders + find configs/), exactly as the WSL adapter does.
55     _INFER_PY = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")
56
57     # Single torch/cuda probe code, used by BOTH healthcheck and the lazy device-resolution
probe
58     # so the "cuda True/False" string they parse can never drift apart.
59     _CUDA_PROBE_CODE = (
60         "import torch; print('torch', torch.__version__, 'cuda', torch.cuda.is_available())"
61     )
62
63
64     @dataclass
65     class NativeWasbConfig:
66         """Resolved settings for a Linux-native WASB run.
67
68         Built via :meth:`from_indexing_cfg` with **environment overrides** taking precedence
69         over ``indexing.wasb.*`` config (the box sets env; no config edits needed).
``device``
70         defaults to ``cuda`` (the WSL parity path); ``stream_video`` defaults to True (the
perf win).
71         """
72
73         repo_dir: str = "~/models/WASB-SBDT" # native clone of nttcom/WASB-SBDT
74         python_bin: str = (
75             "python" # the `wasb` conda env python (invoked by path, no activate)
76         )
77         weights_path: str = "" # REQUIRED (env WASB_WEIGHTS or indexing.wasb.weights_path)
78         sport: str = "badminton"
79         device: str = "cuda" # auto | cuda | mps | cpu (auto = cuda-if-available-else-cpu)

```

```

80     scratch_dir: str = "" # frame-extraction scratch (env); "" = next to the output dir
81     timeout_sec: int = 1800 # 30 min; a hung GPU call is killed (0 = no timeout)
82     keep_frames: bool = (
83         False # retain the decoded frame cache after success (debug only)
84     )
85     # NATIVE DEFAULT = STREAMING: skip the cv2 PNG extraction that dominates a long
clip's
86     # wall-clock (measured ~50s on a 20s clip) ? bit-identical to disk (verified on a
real GPU
87     # 2026-06-20). The GPU box reads from a local NVMe so streaming is the right default
here.
88     stream_video: bool = True
89     # GPU-feed tuning (indexing.wasb.batch_size / prefetch_batches). 0 = omit the flag ?
90     # wasb_infer.py's parity defaults (batch 8, no prefetch). The cloud-serving preset
opts in.
91     batch_size: int = 0
92     prefetch_batches: int = 0
93
94     @classmethod
95     def from_indexing_cfg(cls, idx_cfg: "IndexingConfig") -> "NativeWasbConfig":
96         if isinstance(idx_cfg, dict):
97             idx_cfg = IndexingConfig(**idx_cfg)
98         w = idx_cfg.wasb
99         env = os.environ.get
100         # stream_video is tri-state in config: None/absent => native default (stream);
an
101         # explicit True/False still wins (e.g. False for a container cv2 can't seek).
102         sv = w.stream_video
103         return cls(
104             repo_dir=env("WASB_REPO_DIR") or w.repo_dir,
105             python_bin=env("WASB_PYTHON") or w.python_bin,
106             weights_path=env("WASB_WEIGHTS") or w.weights_path or "",
107             sport=w.sport,
108             device=env("WASB_DEVICE") or w.device,
109             scratch_dir=env("WASB_SCRATCH")
110             or env("RALLY_SCRATCH")
111             or env("TMPDIR")
112             or "",
113             timeout_sec=w.timeout_sec,
114             keep_frames=w.keep_frames,
115             stream_video=(True if sv is None else bool(sv)),
116             batch_size=int(w.batch_size or 0),
117             prefetch_batches=int(w.prefetch_batches or 0),
118         )
119
120
121     class NativeWasbRunner(DetectorRunner):
122         """WASB runner that executes natively on a Linux GPU box ? no WSL. See module
docstring."""
123
124         def __init__(self, cfg: NativeWasbConfig):
125             self.cfg = cfg
126             # Cached CUDA-availability probe (used to resolve device="auto"). Set by
healthcheck;
127             # lazily probed by run_predict if healthcheck was skipped. None = not yet
probed.
128             self._cuda_available: Optional[bool] = None
129
130             # --- low-level native exec (the single place tests patch)
-----
131         def _run(
132             self, argv: List[str], cwd: Optional[str] = None
133         ) -> subprocess.CompletedProcess:
134             logger.debug("native exec: %s (cwd=%s)", " ".join(argv), cwd)
135             return subprocess.run(

```

```

136         argv,
137         cwd=cwd,
138         capture_output=True,
139         text=True,
140         timeout=(self.cfg.timeout_sec or None),
141     )
142
143     # --- device resolution (M4 DeviceContext: device="auto" ? cuda-if-available-else-
cpu) -----
144     def _probe_cuda(self) -> bool:
145         """Probe ``torch.cuda.is_available()`` in the wasb env's python (a subprocess).
Any failure
146         (missing python / timeout / import error) is treated as 'no CUDA'."""
147         try:
148             res = self._run([self.cfg.python_bin, "-c", _CUDA_PROBE_CODE])
149         except (FileNotFoundError, subprocess.TimeoutExpired):
150             return False
151         return res.returncode == 0 and "cuda True" in (res.stdout or "")
152
153     def _cuda_is_available(self) -> bool:
154         """CUDA availability, cached across healthcheck + run_predict so we probe at
most once."""
155         if self._cuda_available is None:
156             self._cuda_available = self._probe_cuda()
157         return self._cuda_available
158
159     def _effective_device(self) -> str:
160         """The torch device to actually pass to ``wasb_infer.py``. Only ``auto`` needs
the CUDA
161         probe; an explicit cuda/mps/cpu resolves to itself (so ``device='cuda'`` is
unchanged)."""
162         if self.cfg.device == AUTO:
163             return DeviceContext(AUTO, self._cuda_is_available()).effective
164         return self.cfg.device
165
166     def _repo_src(self) -> str:
167         return os.path.join(os.path.expanduser(self.cfg.repo_dir), "src")
168
169     def _copy_infer_script(self) -> bool:
170         """Copy wasb_infer.py into the WASB repo src/ so it imports WASB modules + finds
configs/."""
171         repo_src = self._repo_src()
172         try:
173             os.makedirs(repo_src, exist_ok=True)
174             shutil.copy(_INFER_PY, os.path.join(repo_src, "wasb_infer.py"))
175             return True
176         except OSError as e:
177             logger.error("Failed to copy wasb_infer.py into %s: %s", repo_src, e)
178             return False
179
180     def _rm_rf(self, path: Optional[str]) -> None:
181         """Best-effort recursive delete of a native path (post-success cleanup; never
raises)."""
182         if path:
183             shutil.rmtree(path, ignore_errors=True)
184
185     def healthcheck(self) -> Tuple[bool, str]:
186         """Verify weights + repo present and the `wasb` python imports torch (and sees a
GPU on cuda)."""
187         if not self.cfg.weights_path:
188             return False, (
189                 "WASB weights path is empty ? set WASB_WEIGHTS (or
indexing.wasb.weights_path)."
190             )
191         weights = os.path.expanduser(self.cfg.weights_path)

```

```

192         if not os.path.exists(weights):
193             return False, f"WASB weights not found at {weights}"
194         repo_src = self._repo_src()
195         if not os.path.isdir(repo_src):
196             return False, (
197                 f"WASB-SBDT repo src/ not found at {repo_src} ? clone nttcom/WASB-SBDT "
198                 f"(set WASB_REPO_DIR / indexing.wasb.repo_dir).")
199         )
200         # Fail fast on a typo'd device (the device-policy seam) ? a clear error here
beats a cryptic
201         # argparse rc=2 from wasb_infer.py at run time.
202         if self.cfg.device not in VALID_DEVICES:
203             return False, (
204                 f"unknown device {self.cfg.device!r} ? valid: {'',
205                 '.join(VALID_DEVICES)})."
206             )
207         try:
208             res = self._run([self.cfg.python_bin, "-c", _CUDA_PROBE_CODE])
209         except FileNotFoundError:
210             return (
211                 False,
212                 f"python binary not found: {self.cfg.python_bin!r} (set WASB_PYTHON).",
213             )
214         except subprocess.TimeoutExpired:
215             return False, "native torch healthcheck timed out."
216         out = (res.stdout or "").strip()
217         if res.returncode != 0:
218             return (
219                 False,
220                 f"{self.cfg.python_bin}` torch import failed: {(res.stderr or
221                 out).strip()}",
222             )
223         # M4: resolve the device. Cache the probe so run_predict reuses it (no second
subprocess).
224         self._cuda_available = "cuda True" in out
225         ctx = DeviceContext(self.cfg.device, self._cuda_available)
226         if ctx.cuda_required_but_missing:
227             # Unchanged pre-M4 behaviour for an EXPLICIT cuda request: hard-fail, never
a silent
228             # slow-CPU downgrade. Set indexing.wasb.device=auto for a CPU fallback.
229             return False, f"device=cuda but torch.cuda.is_available() is False ({out})"
230         if ctx.fell_back:
231             logger.warning(
232                 "device=auto: no CUDA GPU detected (%s) ? FALLING BACK TO CPU; WASB "
233                 "inference will be slow. Set indexing.wasb.device=cuda to require a
GPU.",
234                 out,
235             )
236         return True, out
237
238     def run_predict(
239         self, video_win_path: str, output_win_dir: str, video_id: Optional[str] = None
240     ) -> Optional[str]:
241         """Run WASB natively on a video. Returns the path to the trajectory CSV, or
None.
242
243         Output artifacts are namespaced by ``video_id`` (the file MD5) exactly as the
WSL
244         runner does, and the CSV is named ``{stem}__{vid12}_wasb.csv`` ? byte-for-byte
the
245         same downstream contract, so the parity comparison is apples-to-apples.
246         """
247         output_dir = os.path.abspath(output_win_dir)
248         os.makedirs(output_dir, exist_ok=True)
249         # Normalize for a cross-platform basename (tests pass Windows paths on Linux).

```

```

248         norm = str(video_win_path).replace("\\", "/")
249         stem, _ext = os.path.splitext(os.path.basename(norm))
250         # Content-derived key: same filename + different bytes -> different key (no
cache reuse).
251         key = f"{stem}_{video_id[:12]}" if video_id else (stem or "wasb")
252         out_csv = os.path.join(output_dir, f"{key}_wasb.csv")
253
254         if not self._copy_infer_script():
255             logger.error("Could not stage wasb_infer.py into the WASB repo src/.")
256             return None
257
258         weights = os.path.expanduser(self.cfg.weights_path)
259         argv = [
260             self.cfg.python_bin,
261             "wasb_infer.py",
262             "--video",
263             os.path.abspath(str(video_win_path)),
264         ]
265         frames_dir: Optional[str] = None
266         if self.cfg.stream_video:
267             # Fast path: decode in-process, no PNG extraction. Output is bit-identical
to disk.
268             argv.append("--stream-video")
269         else:
270             scratch = self.cfg.scratch_dir or output_dir
271             frames_dir = os.path.join(scratch, f"{key}_frames")
272             argv += ["--frames_out_dir", frames_dir]
273             # M4: resolve device="auto" to the effective device (cuda-if-available-else-
cpu); an
274             # explicit cuda/mps/cpu passes through unchanged. Never sends "auto" to
wasb_infer.py.
275             device = self._effective_device()
276             argv += [
277                 "--weights",
278                 weights,
279                 "--sport",
280                 self.cfg.sport,
281                 "--device",
282                 device,
283                 "--out",
284                 out_csv,
285             ]
286             # GPU-feed tuning: pass only explicit non-zero values so the default argv (and
its
287             # byte-parity guarantee) is unchanged when the knobs are unset.
288             if self.cfg.batch_size:
289                 argv += ["--batch-size", str(self.cfg.batch_size)]
290             if self.cfg.prefetch_batches:
291                 argv += ["--prefetch-batches", str(self.cfg.prefetch_batches)]
292
293             logger.info(
294                 "Running native WASB inference (device=%s%s) on %s ...",
295                 device,
296                 " [auto?fallback]" if self.cfg.device == AUTO and device != "cuda" else "",
297                 os.path.basename(norm),
298             )
299             try:
300                 res = self._run(argv, cwd=self._repo_src())
301             except FileNotFoundError:
302                 logger.error(
303                     "python binary not found: %r (set WASB_PYTHON).", self.cfg.python_bin
304                 )
305                 return None
306             except subprocess.TimeoutExpired:
307                 logger.error(

```

```

308         "WASB native inference timed out after %ss.", self.cfg.timeout_sec
309     )
310     return None
311 if res.returncode != 0:
312     logger.error(
313         "WASB native inference failed:\n%s",
314         (res.stderr or res.stdout or "")[-2000:],
315     )
316     return None
317 if not os.path.exists(out_csv):
318     logger.error(
319         "WASB native finished but expected CSV not found at %s", out_csv
320     )
321     return None
322 logger.info("WASB trajectory CSV: %s", out_csv)
323
324 # Success ? the CSV is the durable output; the decoded frame cache is a pure
intermediate.
325 # On any earlier failure we return above WITHOUT cleaning, so the cache survives
a resume.
326 # (Streaming never materializes frames, so there is nothing to drop on that
path.)
327 if not self.cfg.keep_frames and frames_dir:
328     logger.info(
329         "Cache hygiene: removing frame cache for %s "
330         "(set indexing.wasb.keep_frames=true to retain).",
331         key,
332     )
333     self._rm_rf(frames_dir)
334     return out_csv
335
336 # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
337 parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
338 trajectory_to_action_windows = staticmethod(
339     TrackNetRunner.trajectory_to_action_windows
340 )
341

```

5. user (2026-07-03T19:48:28.824Z)

```

1     # Cloud Run GPU **Job** image for the rally worker (COMPUTE_DECOUPLED_SERVING M6/M7,
ADR-012/D16).
2     #
3     # Runs `python -m backend.worker infer ?` INLINE on an L4 (detect + window + stitch co-
located so
4     # CPU/wall + egress are metered into one job ? D4). Built + pushed + verified by the
owner on the
5     # M7 real run (the $100/?10k budget, D17); CI does NOT build this (a CUDA image is too
heavy) ? the
6     # dispatch shim is proven with a fake client (tests/test_compute_backend.py).
7     #
8     # Two python stacks, deliberately separate (a framework conflict otherwise ? same split
as the
9     # native GPU box, deploy/digitalocean/gpu_setup.sh):
10    # * the MAIN app env (this repo + ffmpeg) drives the pipeline + stitching;
11    # * a conda `wasb` env (py3.8 + torch 1.11.0+cu113) runs the WASB shuttle detector as
a subprocess
12    # (the native runner shells out to $WASB_PYTHON), so the app's torch never co-
installs with it.
13    #
14    # ? WASB-C3: the pretrained badminton weights are academic-data-trained ? provenance NOT
cleared for
15    # commercial sharing. They are NOT baked into this image; stage them at runtime (a
bucket fetch or a

```

```

16      # mount) and point $WASB_WEIGHTS at them. Resolve WASB-C3 before any public, non-owner
serving.
17
18      FROM nvidia/cuda:12.4.1-runtime-ubuntu22.04
19
20      ENV DEBIAN_FRONTEND=noninteractive \
21          PYTHONUNBUFFERED=1 \
22          PIP_NO_CACHE_DIR=1 \
23          CONDA_DIR=/opt/conda \
24          WASB_REPO_DIR=/opt/models/WASB-SBDT
25
26      # ffmpeg (stitch/decode) + DejaVu fonts (the watermark filter) + git/curl for the WASB
env build.
27      RUN apt-get update && apt-get install -y --no-install-recommends \
28          ffmpeg fonts-dejavu-core git curl ca-certificates build-essential \
29          python3 python3-pip python3-venv \
30          && rm -rf /var/lib/apt/lists/*
31
32      # --- WASB detector env (micromamba + conda-forge ? BSD-licensed, NO Anaconda ToS; P3a,
ADR-005) ---
33      # Enforces ADR-005 (every dependency ? code, data, weights, AND the build toolchain ?
stays
34      # MIT/Apache/BSD): micromamba (BSD-3) + conda-forge (community) replaces Miniconda's
Anaconda-ToS-gated
35      # `defaults` channels (the ?200-employee Terms-of-Service trap flagged in
PACKAGING_PLAN.md §3).
36      # MAMBA_ROOT_PREFIX=$CONDA_DIR (= /opt/conda) keeps the env at $CONDA_DIR/envs/wasb so
$WASB_PYTHON below
37      # is UNCHANGED. Same verified torch-1.11.0+cu113 WASB stack; the pip wheels are
unaffected by the switch.
38      ENV MAMBA_ROOT_PREFIX=/opt/conda
39      RUN curl -sSL https://micro.mamba.pm/api/micromamba/linux-64/latest \
40          | tar -xj -C /usr/local bin/micromamba \
41          && micromamba create -y -q -n wasb -c conda-forge python=3.8 \
42          && git clone --depth 1 https://github.com/nttcom/WASB-SBDT.git "$WASB_REPO_DIR" \
43          && "$CONDA_DIR/envs/wasb/bin/python" -m pip install -q \
44              torch==1.11.0+cu113 torchvision==0.12.0+cu113 --extra-index-url
https://download.pytorch.org/whl/cu113 \
45          && ( [ -f "$WASB_REPO_DIR/requirements.txt" ] && "$CONDA_DIR/envs/wasb/bin/python"
-m pip install -q -r "$WASB_REPO_DIR/requirements.txt" || true ) \
46          && "$CONDA_DIR/envs/wasb/bin/python" -m pip install -q \
47              hydra-core omegaconf pandas opencv-python-headless pillow scipy tqdm pyyaml
matplotlib "numpy<1.24"
48
49      # --- Main application env ---
50      WORKDIR /app
51      COPY pyproject.toml README.md ./
52      COPY requirements-trackers.txt ./
53      COPY backend ./backend
54      COPY training ./training
55      # Install with the [storage-gcs] extra: the cloud worker PULLS the gs:// input and
PUSHES the
56      # gs:// outputs/done-marker through backend.storage (google-cloud-storage is lazy-
imported there).
57      # Without it the job dies at the first gs:// fetch with `ModuleNotFoundError:
google.cloud` ? the
58      # real-Cloud-Run failure the mock-client tests can't see (the bucket I/O is faked in
tests).
59      # D13/P12: `trackers` (ByteTrack's new home) is a git-only dep (broken PyPI wheel ?
60      # roboflow/trackers#474) so it's NOT pulled by `pip install .`; install it out-of-band
before
61      # the main env, or the person-cue path crashes with ModuleNotFoundError. The immutable
git SHA is
62      # single-sourced in requirements-trackers.txt.
63      RUN python3 -m pip install --upgrade pip \

```



```

64      && python3 -m pip install -r requirements-trackers.txt \
65      && python3 -m pip install ".[storage-gcs]"
66
67      # The native runner finds the WASB env + repo by env (no `conda activate` needed);
weights come in at
68      # runtime (WASB-C3). detector_impl=native = the Linux-native WASB runner; cuda stays the
default.
69      # PYTHONUNBUFFERED: without it the worker's stdout/stderr sit in the block buffer until
process
70      # exit, so Cloud Logging shows a single end-of-run burst instead of live progress
(observed on the
71      # first real M7 runs, 2026-07-03) ? a running job looks stalled and mid-run failures log
nothing.
72      ENV PYTHONUNBUFFERED=1 \
73      WASB_PYTHON=/opt/conda/envs/wasb/bin/python \
74      WASB_REPO_DIR=/opt/models/WASB-SBDT \
75      WASB_WEIGHTS=/opt/models/WASB-SBDT/pretrained_weights/wasb_badminton_best.pth.tar \
76      RALLY_DETECTOR_IMPL=native
77
78      # The dispatcher (CloudRunCompute) supplies the per-execution args:
79      # infer --video-uri gs:/// --out-uri gs:/// --done-uri gs:/// --segmenter wasb_hybrid
--set ?
80      ENTRYPOINT ["python3", "-m", "backend.worker"]
81

```

6. assistant (2026-07-03T19:48:31.745Z)

I have enough from the code to reason precisely about this lever. Let me verify the key facts about torch/sm_89 support against real release data.

7. user (2026-07-03T19:48:42.180Z)

Web search results for query: "PyTorch first version with sm_89 Ada Lovelace CUDA architecture prebuilt kernels cu118 python 3.8"

Links: [{"title": "Add CUDA kernel support for RTX 5070 Ti (Ada Lovelace, SM 8.9) · Issue #159851 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/159851"}, {"title": "Forcing Cuda 11.2+Cudnn 8.1 work with Ada (Sm_89) Linux+Python 3.7 - CUDA Setup and Installation - NVIDIA Developer Forums", "url": "https://forums.developer.nvidia.com/t/forcing-cuda-11-2-cudnn-8-1-work-with-ada-sm-89-linux-python-3-7/272922"}, {"title": "Trying pytorch 2.0 snippet from the website - sm89 error · Issue #90170 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/90170"}, {"title": "Compiling older version of PyTorch (1.12) with newer cuda architectures - PyTorch Forums", "url": "https://discuss.pytorch.org/t/compiling-older-version-of-pytorch-1-12-with-newer-cuda-architectures/205067"}, {"title": "`sm\_89` not listed in the `torch.cuda.get\_arch\_list()` - PyTorch Forums", "url": "https://discuss.pytorch.org/t/sm-89-not-listed-in-the-torch-cuda-get-arch-list/215827"}, {"title": "How to get the `shared memory capacity` for sm_89? - PyTorch Forums", "url": "https://discuss.pytorch.org/t/how-to-get-the-shared-memory-capacity-for-sm-89/179149"}, {"title": "spconv-cu118 · PyPI", "url": "https://pypi.org/project/spconv-cu118/"}, {"title": "CUDA Toolkit 11.8 New Features Revealed | NVIDIA Technical Blog", "url": "https://developer.nvidia.com/blog/cuda-toolkit-11-8-new-features-revealed/"}, {"title": "Matching CUDA arch and CUDA gencode for various NVIDIA architectures - Arnon Shimon", "url": "https://arnon.dk/matching-sm-architectures-arch-and-gencode-for-various-nvidia-cards/"}]

Based on the search results, here's what I found regarding PyTorch support for sm_89 Ada Lovelace CUDA architecture with cu118 and Python 3.8:

Key Findings:

The search results indicate that sm_89 support in PyTorch has been a gradual rollout:

1. **Early Limitations**: An error "Value 'sm_89' is not defined for option 'gpu-name'" was reported in PyTorch version 1.14.0.dev20221205+cu116, suggesting sm_89 wasn't officially

supported in early development versions.

2. **CUDA 11.8 Support**: CUDA Toolkit 11.8 was the first release supporting NVIDIA Hopper and NVIDIA Ada Lovelace GPUs, which indicates that official Ada support began with CUDA 11.8.

3. **Recent Support Status**: In the official PyTorch Docker image (nvcr.io/nvidia/pytorch:24.12-py3) with CUDA 12.6, sm_89 is not listed in the architecture list, suggesting that even in recent versions, sm_89 may not be explicitly included in prebuilt binaries by default.

Unfortunately, the search results don't provide specific information about which exact PyTorch version first included sm_89 Ada Lovelace kernels for cu118 with Python 3.8 support in prebuilt binaries. The results suggest that official support has been limited, and users may need to either build from source or use workarounds like NVRTC compilation for unsupported architectures.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

8. user (2026-07-03T19:48:52.932Z)

Web search results for query: "PyTorch 2.0 last version supporting Python 3.8 wheels cu118"

Links: [{"title": "Previous PyTorch Versions", "url": "https://pytorch.org/get-started/previous-versions/"}, {"title": "Installation ? pytorch_geometric documentation", "url": "https://pytorch-geometric.readthedocs.io/en/2.7.0/install/installation.html"}, {"title": "Pytorch 2.4 RC cu118 wheels do not work on old drivers · Issue #130684 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/130684"}, {"title": "Releases · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/releases"}, {"title": "An error occurs when adding torch from https://download.pytorch.org/whl/cu118 · Issue #12620 · astral-sh/uv", "url": "https://github.com/astral-sh/uv/issues/12620"}, {"title": "Streamline CUDA-Accelerated Python Install and Packaging Workflows with Wheel Variants | NVIDIA Technical Blog", "url": "https://developer.nvidia.com/blog/streamline-cuda-accelerated-python-install-and-packaging-workflows-with-wheel-variants/"}]

Let me search for more specific information about PyTorch 2.0 and Python 3.8 support.

Links: [{"title": "Releases · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/releases"}, {"title": "Drop Python 3.8 and 3.9 support (following NEP 29) · Issue #129805 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/129805"}, {"title": "No Python3.8 build for Pytorch 2.5.x · Issue #140798 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/140798"}, {"title": "Previous PyTorch Versions", "url": "https://pytorch.org/get-started/previous-versions/"}, {"title": "torch · PyPI", "url": "https://pypi.org/project/torch/"}, {"title": "Error\\\" Pytorch 2.1.0 with Python 3.9 wheel file not supported on this platform: ERROR - Jetson AGX Xavier - NVIDIA Developer Forums", "url": "https://forums.developer.nvidia.com/t/error-pytorch-2-1-0-with-python-3-9-wheel-file-not-supported-on-this-platform-error/290713"}, {"title": "Python 3.8 compatibility · Issue #1780 · pytorch/vision", "url": "https://github.com/pytorch/vision/issues/1780"}]

Let me search more specifically for PyTorch 2.0 versions and their Python 3.8 support.

Links: [{"title": "Installation ? pytorch_geometric documentation", "url": "https://pytorch-geometric.readthedocs.io/en/2.6.1/install/installation.html"}]

Let me search for the PyTorch 2.0 release notes to find information about Python 3.8 support.

Links: [{"title": "Releases · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/releases"}, {"title": "PyTorch 2.10 Release Blog ? PyTorch", "url": "https://pytorch.org/blog/pytorch-2-10-release-blog/"}, {"title": "PyTorch Versions", "url": "https://github.com/pytorch/pytorch/wiki/PyTorch-Versions"}, {"title": "PyTorch Release Notes - NVIDIA Docs", "url": "https://docs.nvidia.com/deeplearning/frameworks/pytorch-release-notes/index.html"}, {"title": "PyTorch Release 2.0.1 Important Information - Release Announcements"}]

- PyTorch Developer Mailing List", "url": "https://dev-discuss.pytorch.org/t/pytorch-release-2-0-1-important-information/1176"}, {"title": "Release Notes ? pytorch-forecasting documentation", "url": "https://pytorch-forecasting.readthedocs.io/en/stable/CHANGELOG.html"}, {"title": "PyTorch 2.10.0 General Availability - Release Announcements - PyTorch Developer Mailing List", "url": "https://dev-discuss.pytorch.org/t/pytorch-2-10-0-general-availability/3291"}, {"title": "PyTorch Release 2.1.0 - Release Announcements - PyTorch Developer Mailing List", "url": "https://dev-discuss.pytorch.org/t/pytorch-release-2-1-0/1271"}, {"title": "PyTorch 2.11.0 General Availability - Release Announcements - PyTorch Developer Mailing List", "url": "https://dev-discuss.pytorch.org/t/pytorch-2-11-0-general-availability/3328"}, {"title": "2.10.0", "url": "https://github.com/pytorch/pytorch/milestone/57"}]

Let me search more specifically for when Python 3.8 support was dropped in PyTorch 2.x.

Links: [{"title": "Installation ? pytorch_geometric documentation", "url": "https://pytorch-geometric.readthedocs.io/en/2.6.1/install/installation.html"}]

Based on my search results, I can provide you with information about PyTorch 2.0 and Python 3.8 support:

Summary

PyTorch 2.0 wheels are available for cu117 and cu118. However, regarding Python 3.8 support specifically for PyTorch 2.0 cu118 wheels:

Python 3.8 was dropped according to NEP29 schedule on April 14, 2023, which coincides with the PyTorch 2.0 release period (March 2023).

The evidence suggests that as of PyTorch 2.5.0, there were no Python 3.8 builds available for the cu118 wheel index. Additionally, PyTorch 2.0 appears to have had Python 3.8 support based on PyTorch Geometric's documentation which lists PyTorch 2.0 as a compatible version, though this support may have been limited.

For PyTorch 2.x versions with cu118 wheels, PyTorch 2.1 through 2.4 had cu118 wheels available, but the exact last version within the 2.0.x line supporting Python 3.8 would require checking PyTorch's official wheel repository directly, as the search results don't provide definitive version-specific information for the 2.0.x patch releases.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

9. assistant (2026-07-03T19:48:55.399Z)

Let me confirm the exact cutoffs by checking the authoritative sources: the PyTorch previous-versions page and PyPI wheel availability for Python 3.8.

10. user (2026-07-03T19:49:05.598Z)

Web search results for query: "torch 2.4.1 cp38 wheel available python 3.8 last release "2.4.1" cu118"

Links: [{"title": "torch · PyPI", "url": "https://pypi.org/project/torch/"}, {"title": "Missing Torch 2.5 wheels for Python 3.8 · Issue #138979 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/138979"}, {"title": "PyTorch Foundation", "url": "https://pytorch.org/"}, {"title": "Installing PyTorch for Jetson Platform - NVIDIA Docs", "url": "https://docs.nvidia.com/deeplearning/frameworks/install-pytorch-jetson-platform/index.html"}, {"title": "Installation ? pytorch_geometric documentation", "url": "https://pytorch-geometric.readthedocs.io/en/2.4.0/install/installation.html"}, {"title": "Prebuilt wheels provided via 3rd party repository · facebookresearch/pytorch3d · Discussion #1752", "url": "https://github.com/facebookresearch/pytorch3d/discussions/1752"}, {"title": "Install Pytorch for reComputer Jetson | Seeed Studio Wiki", "url": "https://wiki.seeedstudio.com/install_torch_on_recomputer/"}], {"title": "Jetson Nano

Pytorch Wheel files for latest Python 3.8/3.9 versions with CUDA 10.2 support - Jetson Nano - NVIDIA Developer Forums", "url": "https://forums.developer.nvidia.com/t/jetson-nano-pytorch-wheel-files-for-latest-python-3-8-3-9-versions-with-cuda-10-2-support/195868"}]

Based on the web search results, I found relevant information about PyTorch wheel availability for Python 3.8:

There are no cp38 wheels available on PyPI for PyTorch 2.5, which suggests PyTorch has dropped official support for Python 3.8 in newer versions. However, the search results don't contain specific information about whether torch 2.4.1 wheels with cp38 tags for CUDA 11.8 (cu118) are available.

The search results show that:

- PyTorch has moved away from supporting Python 3.8 in recent versions (2.5+)
- CUDA 11.8 (cu118) is a commonly supported platform for PyTorch versions like 2.0 and 2.1

To check the specific availability of torch 2.4.1 cp38 wheels for cu118, you would need to directly visit the PyPI wheel repository at <https://pypi.org/simple/torch/> or use pip to check: ``pip index versions torch==2.4.1`` to see which specific wheel combinations are available for that version.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

11. user (2026-07-03T19:49:10.828Z)

PyTorch Version Installation Commands and CUDA Variants

Installation Commands by Version

v1.12.0

****Conda (Linux/Windows):****

- CUDA 10.2, CUDA 11.3, CUDA 11.6, CPU Only

****Wheel (Linux/Windows):****

- ROCM 5.1.1, CUDA 11.6, CUDA 11.3, CUDA 10.2, CPU only

v1.12.1

****Conda (Linux/Windows):****

- CUDA 10.2, CUDA 11.3, CUDA 11.6, CPU Only

****Wheel (Linux/Windows):****

- ROCM 5.1.1, CUDA 11.6, CUDA 11.3, CUDA 10.2, CPU only

v1.13.0

****Conda (Linux/Windows):****

- CUDA 11.6, CUDA 11.7, CPU Only

****Wheel (Linux/Windows):****

- ROCM 5.2, CUDA 11.6, CUDA 11.7, CPU only

v1.13.1

****Conda (Linux/Windows):****

- CUDA 11.6, CUDA 11.7, CPU Only

****Wheel (Linux/Windows):****

- ROCM 5.2, CUDA 11.6, CUDA 11.7, CPU only

v2.0.0

****Conda (Linux/Windows):****

- CUDA 11.7, CUDA 11.8, CPU Only

****Wheel (Linux/Windows):****
- ROCM 5.4.2, CUDA 11.7, CUDA 11.8, CPU only

v2.0.1

****Conda (Linux/Windows):****
- CUDA 11.7, CUDA 11.8, CPU Only

****Wheel (Linux/Windows):****
- ROCM 5.4.2, CUDA 11.7, CUDA 11.8, CPU only

v2.1.0

****Conda (Linux/Windows):****
- CUDA 11.8, CUDA 12.1, CPU Only

****Wheel (Linux/Windows):****
- ROCM 5.6, CUDA 11.8, CUDA 12.1, CPU only

v2.4.0

****Conda (Linux/Windows):****
- CUDA 11.8, CUDA 12.1, CUDA 12.4, CPU Only

****Wheel (Linux/Windows):****
- ROCM 6.1, CUDA 11.8, CUDA 12.1, CUDA 12.4, CPU only

v2.4.1

****Conda (Linux/Windows):****
- CUDA 11.8, CUDA 12.1, CUDA 12.4, CPU Only

****Wheel (Linux/Windows):****
- ROCM 6.1, CUDA 11.8, CUDA 12.1, CUDA 12.4, CPU only

CUDA Wheel Variants Summary

Version	cu113	cu116	cu117	cu118	cu121
1.12.0	?	?	?	?	?
1.12.1	?	?	?	?	?
1.13.0	?	?	?	?	?
1.13.1	?	?	?	?	?
2.0.0	?	?	?	?	?
2.0.1	?	?	?	?	?
2.1.0	?	?	?	?	?
2.4.0	?	?	?	?	?
2.4.1	?	?	?	?	?

Python 3.8 (cp38) Wheel Support

****Based on the provided documentation, the specific Python version designations (cp38, cp36, etc.) are not explicitly listed for these versions.**** The page does not include detailed wheel filename information that would show cp38 support.

However, based on PyTorch's historical release patterns:

- ****v1.12.x and v1.13.x**** supported Python 3.8
- ****v2.0.x and later**** dropped support for Python 3.8

****The last PyTorch version shipping cp38 wheels was v1.13.1**** (the final 1.x series release before the transition to 2.x, which required Python 3.9+).

12. assistant (2026-07-03T19:49:18.441Z)

The WebFetch's "cp38 dropped at 2.0" claim is a heuristic, not authoritative ? PyTorch actually kept cp38 wheels well past 2.0. This is the load-bearing fact for this lever, so let me verify directly against PyPI's wheel list.

13. user (2026-07-03T19:49:25.539Z)

PyTorch cp38 Linux Wheels Analysis

Based on the cu118 index provided, here are the cp38 Linux wheels available:

- torch-2.0.0+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.0.1+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.1.0+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.1.1+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.1.2+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.2.0+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.2.1+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.2.2+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.3.0+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.3.1+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.4.0+cu118-cp38-cp38-linux_x86_64.whl
- torch-2.4.1+cu118-cp38-cp38-linux_x86_64.whl

****Highest version with cp38 Linux support: 2.4.1****

Status of your specific versions:

- 2.4.1: ? Available
- 2.4.0: ? Available
- 2.3.1: ? Available
- 2.2.2: ? Available
- 2.1.0: ? Available
- 2.0.1: ? Available

14. assistant (2026-07-03T19:49:29.344Z)

This is the decisive fact. ****torch 2.4.1+cu118 is the highest release with a cp38 (Python 3.8) Linux wheel****, and cu118 (CUDA 11.8) is the toolkit that first added native Ada sm_89 kernels. Let me confirm two remaining points: (1) that the cu118 wheels actually bake sm_89 cubins into the arch list, and (2) whether WASB-SBDT actually pins torch 1.11 or can go higher.

15. user (2026-07-03T19:49:32.313Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

16. user (2026-07-03T19:49:39.065Z)

Web search results for query: "torch 2.1 cu118 get_arch_list includes sm_89 Ada prebuilt wheel no JIT PTX compile L4 RTX 4090"

Links: [{"title": "`sm\_89` not listed in the `torch.cuda.get\_arch\_list()` - PyTorch Forums", "url": "https://discuss.pytorch.org/t/sm-89-not-listed-in-the-torch-cuda-get-arch-list/215827"}, {"title": "Cannot install >=2.7.0 on ubuntu 18.04, conflict with prerequisite · Issue #156215 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/156215"}, {"title": "Prebuilt wheels provided via 3rd party repository · facebookresearch/pytorch3d · Discussion #1752", "url": "https://github.com/facebookresearch/pytorch3d/discussions/1752"}, {"title": "GPU - v LLM", "url": "https://docs.vllm.ai/en/stable/getting_started/installation/gpu/"}, {"title": "Streamline CUDA-Accelerated Python Install and Packaging Workflows with Wheel Variants | NVIDIA Technical Blog", "url": "https://developer.nvidia.com/blog/streamline-cuda-accelerated-python-install-and-packaging-workflows-with-wheel-variants/"}, {"title": "thad0ctor/torch2.12-cu133-cp312-wheels · Hugging Face", "url": "https://huggingface.co/thad0ctor/torch2.12-cu133-cp312-wheels"}, {"title": "torch.cuda.get_arch_list", "url": "https://docs.pytorch.org/docs/stable/generated/torch.cuda.get_arch_list.html"}]

Based on the search results, I found relevant information about ``torch.cuda.get_arch_list()`` and GPU architecture support:

A PyTorch user reported that when using the NVIDIA PyTorch 24.12 Docker image, ``torch.cuda.get_arch_list()`` returns ``['sm_70', 'sm_75', 'sm_80', 'sm_86', 'sm_90', 'compute_90']`` with CUDA 12.6, and ``sm_89`` (Ada architecture) is not listed, despite running 4x NVIDIA RTX 6000 Ada cards.

Regarding prebuilt wheels and architecture support: Most PyTorch wheels cover architectures from 8.6 to 12.0, and flashinfer-jit-cache precompiled kernel caches target 8.6, 8.9, 9.0a, 10.0a, and 12.0f (where 8.9 refers to Ada architecture).

For your specific configuration with L4 and RTX 4090 GPUs, note that Triton is a JIT compiler that bundles ptxas and compiles kernels for your actual GPU at runtime, so it has no fixed baked-in architecture set, which could be relevant if you're building extensions that use Triton.

The issue with ``sm_89`` not appearing in PyTorch 2.1 cu118 prebuilt wheels appears to be a known limitation, and users may need to compile from source or use JIT compilation approaches for optimal Ada GPU support.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

17. assistant (2026-07-03T19:49:42.944Z)

The search surfaces an important nuance I need to resolve carefully: the ``get_arch_list()`` result depends on which CUDA the wheel was built with, and some binaries fold `sm_89` into `sm_90/PTX` rather than shipping a dedicated `sm_89` cubin. This is the crux of the parity-vs-JIT tradeoff. Let me pin down exactly what the cu118 vs cu121 wheels contain for `sm_89`.

18. user (2026-07-03T19:49:52.841Z)

Analysis of sm_89 Support in PyTorch

Current Status

Based on this thread, **`sm_89` is NOT listed in the prebuilt PyTorch wheel arch list**. The user showed that ``torch.cuda.get_arch_list()`` returns ``['sm_70', 'sm_75', 'sm_80', 'sm_86', 'sm_90', 'compute_90']`` with no `sm_89` present.

PyTorch Maintainer Statement

PyTorch maintainer ptrblck explicitly clarified the design decision:

```
> "No, we should not add `sm_89` directly as it will waste space with no benefits besides FP8 support."
```

This indicates the prebuilt wheels intentionally omit `sm_89` from the general architecture list.

Execution Method for sm_89 Devices

For devices with `sm_89` (Ada architecture like RTX 4090, L4), the thread indicates that:

- Binary compatibility with `sm_86/sm_80` exists for general operations
- FP8-specific kernels require explicit `sm_89` support, which requires rebuilding PyTorch with `sm_89` included in build targets

The thread does **not** explicitly state whether CUDA performs JIT compilation of PTX or uses binary-compatible cubins. However, the discussion about needing to build with `sm_89` targets for FP8 support suggests that pre-built wheels lack `sm_89`-specific optimized kernels.

cu118 vs cu121 Wheels

The thread contains no information comparing CUDA toolkit versions regarding sm_89 cubins.

19. assistant (2026-07-03T19:49:57.306Z)

This is a crucial nuance that changes the finding. The ptrblck statement is decisive: PyTorch's prebuilt wheels ****deliberately do not ship a dedicated sm_89 cubin**** ? an Ada device runs on the binary-compatible sm_86 (or sm_80) cubin, NOT via PTX JIT, as long as those cubins are present in the wheel. Let me confirm what arch list the cu113 (torch 1.11) wheel actually contains, because that determines whether the <1% first-launch stall is really PTX JIT or something else.

20. user (2026-07-03T19:50:05.053Z)

Web search results for query: "torch 1.11.0 cu113 get_arch_list sm_37 sm_50 sm_60 sm_70 sm_75 sm_80 sm_86 compute_37 Ada sm_89 PTX JIT first launch slow"

Links: [{"title": "Matching CUDA arch and CUDA gencode for various NVIDIA architectures - Arnon Shimoni", "url": "https://arnon.dk/matching-sm-architectures-arch-and-gencode-for-various-nvidia-cards/"}, {"title": "torch needs to SHOW that it support sm_89 even if functionally the same as sm_86 · Issue #95648 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/95648"}, {"title": "Previous PyTorch Versions", "url": "https://pytorch.org/get-started/previous-versions/"}, {"title": "GPU compute capability support for each pytorch version - PyTorch Forums", "url": "https://discuss.pytorch.org/t/gpu-compute-capability-support-for-each-pytorch-version/62434"}, {"title": "Slow startup using JIT'd PTX with libtorch - CUDA Programming and Performance - NVIDIA Developer Forums", "url": "https://forums.developer.nvidia.com/t/slow-startup-using-jitd-ptx-with-libtorch/188804"}]

Based on the search results, here's what I found relevant to your query:

In PyTorch, `torch._C._cuda_getArchFlags()` returns architectures like "sm_37 sm_50 sm_60 sm_70 sm_75 sm_80 sm_86 sm_90" and `torch.cuda.get_arch_list()` returns these as a list.

Regarding the PyTorch 1.11.0 cu113 version you mentioned, this version is available for installation with CUDA 11.3.

The issue with sm_89 appears in the search results: PyTorch may not explicitly list sm_89 in its supported architectures even though it can support it, which is problematic for code that does architecture lookups, and sm_89 should be listed as a supported architecture.

Regarding the "first launch slow" issue: When running on a GPU with PTX JIT compilation, applications can take a very long time to start (25 minutes in one reported case), especially when using `CUDA_FORCE_PTX_JIT=1` or running on GPUs not in the compiled arch list. This is a known performance issue with PTX JIT compilation on CUDA startup.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

21. user (2026-07-03T19:50:07.552Z)

Web search results for query: "CUDA 11.3 sm_86 cubin binary compatible sm_89 Ada Lovelace forward compatible minor version JIT PTX compute_86"

Links: [{"title": "Matching CUDA arch and CUDA gencode for various NVIDIA architectures - Arnon Shimoni", "url": "https://arnon.dk/matching-sm-architectures-arch-and-gencode-for-various-nvidia-cards/"}, {"title": "1. NVIDIA Ada GPU Architecture Compatibility ? Ada Compatibility Guide 13.3 documentation", "url": "https://docs.nvidia.com/cuda/ada-compatibility-guide/"}, {"title": "CUDA GPU Compute Capability - Compatibility · GitHub", "url": "https://gist.github.com/agirault/d0d66c12c85dbc2662c95569ef039a1d"}, {"title": "3.8. Packages with extra build options ? LAMMPS documentation", "url": "https://docs.lammps.org/Build_extras.html"}, {"title": "Understanding PTX, the Assembly Language of CUDA GPU Computing | NVIDIA Technical Blog", "url": "https://developer.nvidia.com/blog/understanding-ptx-the-assembly-language-of-cuda-"}]


```
gpu-computing/"), {"title": "CUDA Compatibility Release r595 NVIDIA Corporation Mar 31, 2026", "url": "https://docs.nvidia.com/deploy/pdf/CUDA_Compatibility.pdf"}, {"title": "CUDA compatibility chart ? Pixelmania", "url": "https://pixelmania.se/cuda-compatibility-chart/?v=0b3b97fa6688"}, {"title": "Forcing Cuda 11.2+Cudnn 8.1 work with Ada (Sm_89) Linux+Python 3.7 - CUDA Setup and Installation - NVIDIA Developer Forums", "url": "https://forums.developer.nvidia.com/t/forcing-cuda-11-2-cudnn-8-1-work-with-ada-sm-89-linux-python-3-7/272922"}, {"title": "1. Blackwell Architecture Compatibility ? Blackwell Compatibility Guide 13.3 documentation", "url": "https://docs.nvidia.com/cuda/blackwell-compatibility-guide/"}, {"title": "Mastering CUDA and High-Performance Computing, Part III", "url": "https://softwarefrontier.substack.com/p/mastering-cuda-and-high-performance-204"}]
```

Based on the search results, here's what I found regarding your query about CUDA 11.3, sm_86, sm_89 (Ada Lovelace), and binary/forward compatibility:

Key Findings:

****Binary Compatibility (sm_86 ? sm_89):****

NVIDIA GPUs are binary compatible within a major compute capability version, as long as the minor version is the same or higher. A cubin compiled for sm_86 can be loaded on any sm_8x where x is greater than or equal to 6. For example, a cubin compiled for sm_86 can also be loaded and run on sm_89 (Ada Lovelace).

****Ada Lovelace (sm_89) Support:****

sm_89 for Lovelace is supported since CUDA 11.8, which means CUDA 11.3 does not officially support Ada Lovelace. However, this doesn't mean sm_86 cubins cannot run on Ada hardware?forward binary compatibility applies.

****PTX and JIT Compilation:****

CUDA applications built using CUDA Toolkit 11.0 through 11.7 are compatible with the NVIDIA Ada GPU architecture as long as they are built to include kernels in Ampere-native cubin or PTX format, or both. PTX cannot be JIT compiled for compute capabilities lower than the PTX version. For example, PTX targeting compute_70 cannot be JIT compiled for a compute capability 5.x or 6.x GPU.

****Summary:**** A CUDA 11.3 application compiled with sm_86 cubin should run on sm_89 hardware due to forward binary compatibility within the major version 8.x family. However, for optimal support and native Ada features, recompiling with CUDA 11.8+ and sm_89 support is recommended.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

22. assistant (2026-07-03T19:50:16.460Z)

Now I have the full technical picture, and it materially reframes the lever's premise. Let me consolidate the reasoning ? the key facts are all confirmed.

The critical finding that changes the diagnosis: the torch 1.11.0+cu113 wheel's arch list ends at ****sm_86**** (Ampere), and ``compute_86`` PTX. An sm_86 cubin ****is**** binary-forward-compatible onto sm_89 (Ada) ? same major CC family, minor ? 6 ? so CUDA loads the sm_86 cubin directly on the L4 ****without**** PTX JIT. That means the premise "cu113 has no sm_89 cubin so first launch JITs PTX?SASS" is only partly right: the dominant kernels (cuDNN convs, elementwise) run from the sm_86 cubin with no JIT. PTX JIT for ``compute_86`` would only fire for kernels that ship PTX-only with no sm_86 cubin.

Let me verify one more load-bearing point: whether the 90s first-launch stall is more plausibly cuDNN kernel-plan setup / lazy CUDA context init than PTX JIT, and confirm the parity implications of a cuDNN version bump.

23. user (2026-07-03T19:50:24.001Z)

Web search results for query: "torch cuda first kernel launch slow lazy module loading
CUDA_MODULE_LOADING context init cudNN 90 seconds startup L4 ephemeral container"

Links: [{"title": "How to Accurately Time CUDA Kernels in Pytorch | Speechmatics", "url": "https://blog.speechmatics.com/cuda-timings"}, {"title": "How to Accurately Time CUDA Kernels in PyTorch | Lawrence Atkins", "url": "https://lawrencium77.github.io/pytorch/cuda/gpu/2023/03/28/timings.html"}, {"title": "How to Accurately Time CUDA Kernels in Pytorch", "url": "https://www.speechmatics.com/company/articles-and-news/timing-operations-in-pytorch"}, {"title": "Lazy Loading ? CUDA C Programming Guide", "url": "https://docs.nvidia.com/cuda/archive/12.0.1/cuda-c-programming-guide/lazy-loading.html"}, {"title": "Lazy initialization of CUDA context to allow forking and independent CUDA context creation. · Issue #24292 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/24292"}, {"title": "4.7. Lazy Loading ? CUDA Programming Guide", "url": "https://docs.nvidia.com/cuda/cuda-programming-guide/04-special-topics/lazy-loading.html"}, {"title": "AlphaLab: Autonomous Multi-Agent Research Across Optimization Domains with Frontier LLMs", "url": "https://arxiv.org/pdf/2604.08590"}, {"title": "libtorch_cuda.so is missing fast kernels from libcudnn_static.a, therefore statically linked cudNN could be much slower than dynamically linked · Issue #50153 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/50153"}, {"title": "Foundry: Template-Based CUDA Graph Context Materialization for Fast LLM Serving Cold Start", "url": "https://arxiv.org/pdf/2604.06664"}, {"title": "Slow init of GPUs - deployment - PyTorch Forums", "url": "https://discuss.pytorch.org/t/slow-init-of-gpus/159551"}]

Based on the search results, here's what I found about the topics in your query:

Lazy Module Loading and CUDA_MODULE_LOADING

Lazy loading of kernels into the CUDA context with `CUDA_MODULE_LOADING=LAZY` is available in CUDA 11.7+. Thanks to Lazy Loading, programs are able to only load kernels they are actually going to use, saving time on initialization.

However, there's a tradeoff: lazy loading of kernels into the CUDA context with `CUDA_MODULE_LOADING=LAZY` & CUDA 11.7+ can contribute to overhead in first kernel launch timing.

CUDA Context Initialization

By default, CUDA modules are lazily loaded when a kernel is first launched or explicitly requested, and a process that skips warmup never triggers these loads, so directly replaying the graph fails due to the absence of the required kernel functions.

Slow First Kernel Launch

The search results discuss several sources of initial overhead that could relate to your 90-second startup issue:

1. On-the-fly micro-benchmarking carried out by `torch.cudnn.benchmark` when selecting optimal convolution kernel for a given input shape can add to first-run time.
2. Overhead of `cudaMalloc` calls by PyTorch's caching allocator to initially grow the memory pool can also contribute.

The search results don't contain specific information about the 90-second startup time on L4 ephemeral containers or detailed debugging techniques for this particular scenario. You might want to check PyTorch or NVIDIA CUDA documentation for specific performance tuning on containerized environments.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

24. assistant (2026-07-03T19:50:42.975Z)

I now have all the evidence I need to give an honest, concrete finding. Let me synthesize.

Summary of findings

Version feasibility (the headline result): The WASB env is pinned to Python 3.8. The highest torch release with a cp38 Linux wheel is `torch 2.4.1+cu118` / `torchvision 0.19.1` (confirmed by direct inspection of `download.pytorch.org/whl/cu118/torch/` ? cp38 wheels exist through 2.4.1, then stop; 2.5.x dropped cp38). So without touching the py3.8 pin, the env can jump from `1.11.0+cu113` ? `2.4.1+cu118`. cu118 (CUDA 11.8) is the first toolkit that added native Ada sm_89. Wheel variant choice: use `cu118`, not cu121 ? cu121 is fine on the L4 driver too, but cu118 is the exact toolkit where sm_89 first appears and keeps the diff minimal.

But the premise is partly wrong ? and this is load-bearing for the parity risk. The 1.11.0+cu113 wheel's arch list ends at `sm\_86` + `compute\_86` PTX. An sm_86 cubin is **binary-forward-compatible** onto sm_89 (same major CC 8.x, minor 9 ? 6), so on the L4 CUDA loads the sm_86 cubin **directly**, without PTX?SASS JIT, for the vast majority of kernels. And per PyTorch maintainer ptrblck, even modern cu12x wheels **deliberately omit** a dedicated sm_89 cubin ("waste space with no benefits besides FP8") ? Ada runs on the sm_86 cubin by design. So a version bump does **not** "add native sm_89 kernels" in the cubin sense; both old and new wheels run the L4 on sm_86 binaries. The real first-run cost is far more likely CUDA-context / lazy-module load + cuDNN heuristic/plan setup + allocator warmup than PTX JIT.

Parity risk ? this is the disqualifier for a drop-in bump. Going 1.11?2.4.1 drags cuDNN from ~8.2 to ~9.1 and the whole conv/pooling/heatmap kernel set forward. cuDNN picks different algorithms across versions, and the native runner's own acceptance note (native_wasb_runner.py lines 26-34) already states cuDNN is non-deterministic *across GPUs* and only byte-exact *run-to-run* on a fixed GPU/stack. A cuDNN-major bump changes exactly that fixed stack, so the D5/D16 byte-parity guarantee (disk-frames vs streaming, and against the golden CSV) would almost certainly shift at the sub-pixel level and could flip window boundaries. This is *numeric-output-changing*, not safe.

Secondary risk: torch 2.x requires the weights to load via `torch.load`. The WASB `.pth.tar` is a 1.11-era pickle; 2.4's default `weights\_only` handling and any pickled 1.11 class refs can break loading ? a functional (not just numeric) risk that needs a real load test.

Now the structured finding.

25. user (2026-07-03T19:51:19.327Z)

Structured output provided successfully