

claude-windows / 6e9d0958-a2ec-4a8e-bb2d-6cce9833c527

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\6e9d0958-a2ec-4a8e-bb2d-6cce9833c527\subagents\workflows\wf_8d68d336-288\agent-acb1882e4b4f3d2b9.jsonl`
- SHA-256: `5163e1f806fb1de5c161812d713e49c0fb27657c1171be06960011604b6e6e7b`
- Source modified: `2026-07-03T19:50:42+00:00`
- Imported at: `2026-07-05T16:48:26+00:00`
- project: `wf_8d68d336-288`
- session_id: `6e9d0958-a2ec-4a8e-bb2d-6cce9833c527`
```

Transcript

1. user (2026-07-03T19:48:21.204Z)

Repo (my in-progress optimization branch): C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree

Key files to Read:

- backend/pipeline/detectors/wasb_infer.py ? the inference loop (DataLoader ? detector.run_tensor ? tracker). Contains the streaming SequentialFrameStore, the _PrefetchIter I just added, build_windows, and the tracker re-run loop.
- backend/pipeline/detectors/native_wasb_runner.py ? builds the wasb_infer.py argv on the GPU box (native runner). Has device/stream_video/batch_size/prefetch_batches config.
- backend/config/models.py (WasbIndexConfig) and backend/config/presets.py (cloud-serving preset).
- deploy/cloudrun/Dockerfile ? the WASB conda env is torch 1.11.0+cu113 (py3.8), a SEPARATE env from the app's torch. The detector shells out to it as a subprocess.

The WASB model itself lives at /opt/models/WASB-SBDT (nttcom/WASB-SBDT) ? you can't read it, but reason from how wasb_infer.py calls build_detector(cfg) / detector.run_tensor(imgs, trans) and build_tracker(cfg).

MEASURED BASELINE (first real L4 cloud run, 2026-07-03): a ~7.5-min 60fps badminton clip took Phase-2 (WASB detection) 727.9s (~0.6x realtime); GPU utilization sat <1% the whole time. Hardware: NVIDIA L4 (Ada sm_89, 23GB), 8 vCPU, 32GB. Streaming mode forces DataLoader num_workers=0 (single-process), so CPU batch assembly (cv2 decode + PIL + ToTensor) serializes with GPU inference.

CRITICAL CONSTRAINT ? the D5/D16 byte-parity guarantee: the repo REQUIRES that the rally/window OUTPUT (the candidate windows and trajectory CSV) stay bit-identical / deterministic across runs and across the disk-frames vs streaming paths. This is asserted by a real parity test. An optimization that changes detector NUMERICS (FP16, TF32, cudnn.benchmark algo selection, NVDEC vs CPU decode) risks changing the output and breaking that guarantee. Batch size and prefetch do NOT change per-frame numerics (same forward, just grouped/overlapped) so they are parity-safe. Already implemented on this branch: indexing.wasb.batch_size + prefetch_batches knobs, the cloud-serving preset opts into batch 64 + prefetch 4.

You are a CUDA/PyTorch performance engineer. Investigate ONE optimization lever in depth by READING the actual code. Lever "tracker-postproc": Investigate the tracker re-run loop at the end of wasb_infer.py run(): after detection it iterates sorted(det_by_fid) calling tracker.update() per frame on CPU. For a long video (tens of thousands of frames) is this a serial CPU bottleneck that runs AFTER the GPU is idle? Measure its likely cost vs the detection pass. Is it parallelizable or vectorizable? Also check the per-window Python loop that assembles frames_payload and the JSONL cache flush ? any O(n^2) or per-frame Python overhead that dominates on long clips?

Read the named files, reason from the real code, and return the structured finding. Be concrete

about code_locations (file:function). Be honest about parity_risk ? the byte-parity guarantee is load-bearing.

2. assistant (2026-07-03T19:48:24.408Z)

I'll investigate the tracker-postproc lever by reading the actual code. Let me start with the key file.

3. user (2026-07-03T19:48:27.172Z)

```
1      """WASB shuttle-trajectory inference on an arbitrary frame sequence (runs in WSL).
2
3      WASB-SBDT only ships *dataset* evaluation; this is a thin wrapper that reuses its
4      own building blocks (ImageDataset transform + detector + online tracker) to run on
5      any folder of extracted frames and emit a trajectory CSV in the exact shape the
6      indexer's `tracknet_runner.parse_trajectory_csv` already consumes:
7
8          Frame,Visibility,X,Y
9
10     It must run INSIDE the WSL `wasb` conda env, from the WASB-SBDT `src/` dir (it
11     imports WASB's `detectors`/`trackers`/`dataloaders`). The Windows-side adapter
12     (`wasb_runner.py`) stages this file + the frames into WSL and invokes it.
13
14     Resumable across reboots
15     -----
16     The GPU detector pass (one forward per sliding window ? ~21k for a 6-min clip) is
17     the slow, expensive part; the CPU tracker pass that turns detections into a
18     trajectory is cheap and *stateful* (it must see every frame in order, so it can't
19     resume mid-stream). We exploit that split:
20
21     * The detector pass writes its per-window raw outputs incrementally to
22       ``<cache>/det_raw.jsonl`` (flushed+fsync'd every ``--flush-every`` windows) and
23       records progress in ``<cache>/manifest.json`` (atomic write). A reboot loses at
24       most ``--flush-every`` windows of GPU work instead of the whole pass.
25     * On restart we replay the cached windows, skip the DataLoader ahead to the first
26       un-cached window, and continue. Once every window is cached, the detector pass
27       is skipped entirely and we just re-run the (cheap, deterministic) tracker over
28       the full ordered cache -> trajectory CSV. So a lost trajectory CSV costs seconds,
29       not the GPU hour.
30
31     The cache is keyed by the output path (``<out>_wasbcache/`` by default), lives next
32     to the output (NOT in %TEMP%), and is invalidated automatically if the frames dir,
33     weights, sport, window size, or frame count change.
34
35     Usage (in WSL, from ~/models/WASB-SBDT/src):
36         python wasb_infer.py --frames_dir /path/frames --weights
37         ../pretrained_weights/wasb_badminton_best.pth.tar \
38         --out /path/traj.csv [--sport badminton] [--limit N] [--cache-dir DIR] [--flush-
39         every 1000] [--fresh]
40
41     """
42
43     import argparse
44     import csv
45     import glob
46     import json
47     import logging
48     import os
49     import os.path as osp
50     from collections import defaultdict
51     from contextlib import contextmanager
52
53     logger = logging.getLogger(__name__)
```

```

54     _MANIFEST_FILE = "manifest.json"
55
56
57     def _build_cfg(weights: str, sport: str, device: str = "cuda"):
58         """Compose the WASB eval config via Hydra, overriding model/weights/device.
59
60         ``device`` is ``cuda`` (default ? the historical WSL behaviour, byte-parity
61         preserved),
62         ``mps`` (Apple) or ``cpu``. Only the single-GPU CUDA path requests
63         ``runner.gpus=[0]``;
64         cpu/mps run with no GPU list so the detector places tensors on ``device``.
65         """
66         from hydra import compose, initialize
67
68         gpus = "[0]" if device == "cuda" else "[]" # single GPU on cuda; none for cpu/mps
69         with initialize(config_path="configs", version_base=None):
70             cfg = compose(
71                 config_name="eval",
72                 overrides=[
73                     f"dataset={sport}",
74                     "model=wasb",
75                     f"detector.model_path={weights}",
76                     f"runner.device={device}",
77                     f"runner.gpus={gpus}", # default config assumes multiple GPUs; pin to
78                     this box
79                 ],
80             )
81             return cfg
82
83     def _frame_id(path: str) -> int:
84         """Best-effort integer frame id from a frame filename (e.g. frame_000180.png ->
85         180)."""
86         stem = osp.splitext(osp.basename(path))[0]
87         digits = "".join(ch for ch in stem if ch.isdigit())
88         return int(digits) if digits else -1
89
90     # ----- #
91     # Resumable detector cache (pure-Python, no torch/numpy ? unit-testable)
92     # ----- #
93     def _cache_paths(cache_dir: str):
94         return osp.join(cache_dir, _DET_FILE), osp.join(cache_dir, _MANIFEST_FILE)
95
96     def default_cache_dir(out_csv: str) -> str:
97         """Stable cache dir next to the trajectory output (survives reboots; not %TEMP%)."""
98         d = osp.dirname(osp.abspath(out_csv))
99         stem = osp.splitext(osp.basename(out_csv))[0]
100         return osp.join(d, stem + "_wasbcache")
101
102     def _atomic_write_text(path: str, text: str) -> None:
103         """Write then rename so a crash mid-write can't leave a half-written file."""
104         tmp = path + ".tmp"
105         with open(tmp, "w") as f:
106             f.write(text)
107             f.flush()
108             os.fsync(f.fileno())
109         os.replace(tmp, path)
110
111     def _det_plain(det) -> list:
112         """A detector candidate dict {'xy': np.array([x,y]), 'score', 'scale'} -> JSON-able
113         list."""

```

```

114     xy = det["xy"]
115     return [float(xy[0]), float(xy[1]), float(det["score"]), int(det["scale"])]
116
117
118     def _dets_to_tracker(plain_list):
119         """Inverse of _det_plain: rebuild the dicts the tracker expects (xy MUST be a numpy
array,
120         the tracker does vector math / np.linalg.norm on it)."""
121         import numpy as np
122
123         return [
124             {"xy": np.array([p[0], p[1]], dtype=float), "score": p[2], "scale": p[3]}
125             for p in plain_list
126         ]
127
128
129     def write_manifest(cache_dir: str, manifest: dict) -> None:
130         _, mpath = _cache_paths(cache_dir)
131         _atomic_write_text(mpath, json.dumps(manifest, indent=2))
132
133
134     def read_manifest(cache_dir: str):
135         _, mpath = _cache_paths(cache_dir)
136         if not osp.exists(mpath):
137             return None
138         try:
139             with open(mpath) as f:
140                 return json.load(f)
141         except (json.JSONDecodeError, OSError):
142             return None
143
144
145     def manifest_compatible(
146         manifest: dict,
147         *,
148         frames_dir: str,
149         weights: str,
150         sport: str,
151         frames_in: int,
152         frames_total: int,
153         device: str = "cuda",
154     ) -> bool:
155         """A cache is reusable only if the run that produced it matches this run's inputs.
156
157         ``device`` defaults to ``cuda`` so caches written before device-keying (no
158         ``device``
159         field) stay compatible with a CUDA run ? but a cpu/mps run will NOT reuse a cuda
160         cache
161         (detections can differ numerically across devices), and vice-versa.
162
163         ``frames_dir`` is the **already-canonical cache key** the caller stored in the
164         manifest
165         (``run()`` passes its ``cache_key`` here, and writes the same value into the
166         manifest):
167         ``osp.abspath(frames_dir)`` for the frames-on-disk path, or the synthetic
168         ``video:<abspath>``
169         key for the streaming path. We therefore compare it VERBATIM ? do NOT re-``abspath``
170         it here:
171         ``osp.abspath("video:/x.mp4")`` is not recognised as absolute and gets the CWD
172         prepended,
173         so a streaming cache would never match itself and every re-run would recompute
174         (issue #348).
175         """
176         if not manifest or manifest.get("version") != CACHE_VERSION:
177             return False

```

```

170     return (
171         manifest.get("frames_dir") == frames_dir
172         and manifest.get("weights") == weights
173         and manifest.get("sport") == sport
174         and manifest.get("frames_in") == frames_in
175         and manifest.get("frames_total") == frames_total
176         and manifest.get("device", "cuda") == device
177     )
178
179
180 def reset_cache(cache_dir: str) -> None:
181     """Drop any existing cache so the next run starts clean."""
182     det_jsonl, mpath = _cache_paths(cache_dir)
183     for p in (det_jsonl, mpath):
184         if osp.exists(p):
185             os.remove(p)
186
187
188 def load_and_clean_cache(cache_dir: str):
189     """Read the longest *contiguous* (w == line index) valid prefix of det_raw.jsonl,
190     rebuild the per-frame detection accumulation, and rewrite the file to exactly that
191     prefix so a trailing half-written line from a crash can't corrupt future appends.
192
193     Returns (windows_done, det_by_fid) where det_by_fid maps frame_id -> list of
194     [x, y, score, scale] candidates (accumulated across every window the frame is in,
195     in window order ? identical to a non-resumed run).
196     """
197     det_jsonl, _ = _cache_paths(cache_dir)
198     det_by_fid = defaultdict(list)
199     valid: list = []
200     if osp.exists(det_jsonl):
201         with open(det_jsonl, "r") as f:
202             for line in f:
203                 s = line.strip()
204                 if not s:
205                     continue
206                 try:
207                     obj = json.loads(s)
208                 except json.JSONDecodeError:
209                     break # truncated trailing line (crash mid-flush)
210                 if obj.get("w") != len(valid): # non-contiguous -> stop at the gap
211                     break
212                 valid.append(s)
213                 for fid, plain in obj["f"]:
214                     det_by_fid[int(fid)].extend([list(p) for p in plain])
215                 # Guarantee a clean append boundary by rewriting only the good prefix.
216                 _atomic_write_text(det_jsonl, ("\n".join(valid) + "\n") if valid else "")
217     return len(valid), det_by_fid
218
219
220 def _append_windows(fh, objs) -> None:
221     """Append window records as JSONL and durably flush (bounded loss on crash)."""
222     for o in objs:
223         fh.write(json.dumps(o, separators=(",", ":")) + "\n")
224     fh.flush()
225     os.fsync(fh.fileno())
226
227
228 def _atomic_write_trajectory(out_csv: str, rows) -> None:
229     os.makedirs(osp.dirname(osp.abspath(out_csv)), exist_ok=True)
230     tmp = out_csv + ".tmp"
231     with open(tmp, "w", newline="") as f:
232         w = csv.writer(f)
233         w.writerow(["Frame", "Visibility", "X", "Y"])
234         w.writerows(rows)

```

```

235         f.flush()
236         os.fsync(f.fileno())
237     os.replace(tmp, out_csv)
238
239
240     # ----- #
241     # Frame addressing + windowing (pure; no torch/cv2 ? unit-testable)
242     # ----- #
243     _STREAM_FRAME_FMT = "{:08d}.png"
244
245
246     def synthetic_frame_path(index: int) -> str:
247         """Stable, index-encoding frame name used by the streaming path.
248
249         It mirrors the on-disk names ``extract_frames`` writes (``{i:08d}.png``) so the
250         downstream integer-id parser (``_frame_id``) yields the SAME frame id whether the
251         frames came from disk or from the in-memory stream. The streaming image loader
252         inverts this name back to an index to fetch the decoded frame.
253         """
254         return _STREAM_FRAME_FMT.format(int(index))
255
256
257     def build_windows(frames: list, frames_in: int):
258         """Sliding windows of ``frames_in`` consecutive frame *paths* (the unit of GPU work
259         + checkpoint). Pure list math, identical for the disk and streaming paths.
260
261         Returns a list of windows, each a list of ``frames_in`` consecutive paths
262         (``[frames[i:i+frames_in] for i in range(len(frames) - frames_in + 1)]``). The
263         caller wraps each window into the ``{"frames", "annos"}`` sample shape WASB's
264         ``ImageDataset`` expects; keeping that out of here stays torch-free for unit tests.
265         """
266         return [frames[i : i + frames_in] for i in range(len(frames) - frames_in + 1)]
267
268
269     class SequentialFrameStore:
270         """In-memory sliding buffer over a forward-only frame reader, sized for the
271         detector's sliding-window access pattern (peak O(window) frames, not O(video)).
272
273         The detector DataLoader runs with ``shuffle=False`` and (in streaming mode)
274         ``num_workers=0``. ``ImageDataset.__getitem__(i)`` reads the frames of window ``i``
275         in order ? ``i, i+1, ..., i+window-1`` ? and samples are requested ``i = start,
276         start+1, ...``. So the global read sequence dips back by ``window-1`` at each window
277         boundary (``0,1,2, 1,2,3, 2,3,4, ...`` for ``window=3``); the highest index seen so
278         far never decreases. We keep only frames at or above ``max_seen - (window-1)`` (the
279         earliest index any not-yet-finished window can still ask for) and decode each source
280         frame exactly once via the forward-only ``reader(index)``.
281         """
282
283         def __init__(self, reader, total: int, window: int = 1, start_index: int = 0):
284             self._reader = reader
285             self._total = int(total)
286             self._window = max(1, int(window))
287             self._buf: dict = {}
288
289             # On a resumed run the detector's first needed frame is `start_index`; begin
290             # pulling there so the reader can skip (seek past) the already-cached prefix
291             # instead of decoding 0..start_index-1 just to throw them away.
292             self._next = int(start_index) # next index not yet pulled from the reader
293             self._max_seen = int(start_index) - 1
294             self._floor = int(start_index) # lowest index we still keep (never evict below)
295
296         def get(self, index: int):
297             index = int(index)
298             if index < self._floor:
299                 raise KeyError(
300                     f"frame {index} already evicted (below floor {self._floor}); the access

```

```

"
300         f"pattern must stay within a window of the highest frame seen"
301     )
302     if index >= self._total:
303         raise KeyError(f"frame {index} out of range (total {self._total})")
304     # Pull forward from the reader until the requested index is buffered.
305     while self._next <= index:
306         self._buf[self._next] = self._reader(self._next)
307         self._next += 1
308     if index > self._max_seen:
309         self._max_seen = index
310     # Evict frames older than the earliest a still-running window could re-request:
311     # once we've seen index `m`, no future window starts before `m - (window-1)`.
312     new_floor = max(self._floor, self._max_seen - (self._window - 1))
313     for k in range(self._floor, new_floor):
314         self._buf.pop(k, None)
315     self._floor = new_floor
316     return self._buf[index]
317
318
319 def _index_from_synthetic_path(path: str) -> int:
320     """Invert ``synthetic_frame_path``: a frame path -> its integer source index.
321
322     Reuses ``_frame_id`` (digits-only parse) so the index and the cached frame-id stay
323     in lockstep with the on-disk naming scheme.
324     """
325     return _frame_id(path)
326
327
328 def _bgr_to_pil_rgb(frame_bgr):
329     """Convert an in-memory BGR uint8 frame (as ``cv2.VideoCapture`` yields) to the
330     EXACT
331     PIL RGB image WASB's disk path produces.
332
333     The disk path is ``frame_bgr -> cv2.imwrite(PNG) ->
334     Image.open(PNG).convert('RGB')``;
335     because PNG is lossless that round-trip equals ``cv2.cvtColor(frame_bgr,
336     COLOR_BGR2RGB)``
337     bit-for-bit (verified empirically, max|diff|=0). Returning that here makes the
338     streamed
339     frame indistinguishable from the on-disk one to everything downstream.
340     """
341     import cv2
342     from PIL import Image
343
344     return Image.fromarray(cv2.cvtColor(frame_bgr, cv2.COLOR_BGR2RGB))
345
346
347 def _make_streamed_read_image(store, real_read_image):
348     """Build the ``read_image`` replacement used during a streaming detector pass.
349
350     A synthetic frame path -> the in-memory decoded frame for its index (as a PIL RGB
351     image,
352     matching ``read_image``'s real return type); any path that doesn't parse to an index
353     falls
354     through to ``real_read_image``. Pure ? imports no WASB module ? so it is unit-
355     testable
356     without the WSL-only ``dataloaders`` package.
357     """
358
359     def _read_image(path, *args, **kwargs):
360         idx = _index_from_synthetic_path(path)
361         if idx < 0:
362             return real_read_image(path, *args, **kwargs)
363         return _bgr_to_pil_rgb(store.get(idx))

```

```

357
358     return _read_image
359
360
361     @contextmanager
362     def _streaming_read_image_patch(
363         frame_loader, total: int, window: int = 1, start_index: int = 0
364     ):
365         """Scope a shim over WASB's ``read_image`` so ``ImageDataset`` is fed in-memory
366         decoded
367         frames during the detector pass, byte-for-byte as it would over real PNGs.
368
369         WASB's ``ImageDataset.__getitem__`` reads each frame via ``read_image(path)`` ? NOT
370         ``cv2.imread``. ``read_image`` (``utils.utils``) does
371         ``Image.open(path).convert('RGB')``,
372         returning a PIL **RGB** image, after an ``osp.exists(path)`` guard. We therefore
373         feed it a
374         PIL RGB image built from the streamed BGR frame (see ``_bgr_to_pil_rgb``), which is
375         byte-identical to the disk round-trip, and the shim never touches the filesystem so
376         the
377         ``osp.exists`` guard is sidestepped for synthetic stream paths.
378
379         We patch ``dataloaders.dataset_loader.read_image`` ? the binding the consumer
380         actually
381         calls (``dataset_loader`` does ``from utils import read_image``, so the name lives
382         in that
383         module's namespace; patching ``utils.read_image`` would NOT rebind the already-
384         imported
385         reference).
386
387         ``frame_loader(index) -> BGR uint8 ndarray`` is a forward-only sequential decoder
388         wrapped
389         in a ``SequentialFrameStore`` (sliding buffer sized to ``window`` = ``frames_in``);
390         on a
391         resume we start at ``start_index`` so the decoder seeks past already-cached windows.
392         The
393         original reader is restored on exit.
394         """
395         from dataloaders import dataset_loader as _dl
396
397         store = SequentialFrameStore(
398             frame_loader, total, window=window, start_index=start_index
399         )
400         real_read_image = _dl.read_image
401         # Rebind read_image in the consuming module for the duration of the detector pass;
402         # restore on exit. (Plain assignment, not setattr-with-constant, to stay ruff
403         B010-clean.)
404         _dl.read_image = _make_streamed_read_image(store, real_read_image) # type:
405         ignore[assignment]
406         try:
407             yield store
408         finally:
409             _dl.read_image = real_read_image # type: ignore[assignment]
410
411
412     # ----- #
413     # Inference
414     # ----- #
415     class _PrefetchIter:
416         """Bounded producer-thread wrapper over a batch iterator (the num_workers=0
417         DataLoader).
418
419         In streaming mode the DataLoader must stay single-process (the in-memory frame store
420         +
421         the cv2.imread shim live in this process and cannot be pickled to workers), so the

```



```

408 CPU-side batch assembly (decode + PIL + ToTensor) SERIALIZES with GPU inference ?
409 measured on the first real L4 serving run (2026-07-03) this left the GPU <1%
utilized.
410 A single producer THREAD fixes the overlap without multiprocessing: cv2/PIL/numpy
411 release the GIL for the heavy parts, the store's forward-only access pattern is
412 preserved (one consumer thread iterating sequentially), and the bounded queue caps
413 memory at ``depth`` batches. A producer exception is re-raised in the consumer."""
414
415 _END = object()
416
417 def __init__(self, src, depth: int = 2):
418     import queue
419     import threading
420
421     self._q: "queue.Queue" = queue.Queue(maxsize=max(1, int(depth)))
422     self._exc: "BaseException | None" = None
423
424     def _fill() -> None:
425         try:
426             for item in src:
427                 self._q.put(item)
428             except BaseException as e: # noqa: BLE001 - surfaced in __next__
429                 self._exc = e
430             finally:
431                 self._q.put(self._END)
432
433     self._t = threading.Thread(target=_fill, name="wasb-prefetch", daemon=True)
434     self._t.start()
435
436 def __iter__(self):
437     return self
438
439 def __next__(self):
440     item = self._q.get()
441     if item is self._END:
442         if self._exc is not None:
443             raise self._exc
444         raise StopIteration
445     return item
446
447
448 def run(
449     frames,
450     weights: str,
451     out_csv: str,
452     sport: str = "badminton",
453     limit: int = 0,
454     batch_size: int = 8,
455     num_workers: int = 4,
456     log_every_batches: int = 50,
457     cache_dir: "str | None" = None,
458     flush_every: int = 1000,
459     fresh: bool = False,
460     frame_loader=None,
461     source_key: "str | None" = None,
462     device: str = "cuda",
463     prefetch_batches: int = 0,
464 ) -> str:
465     """Run WASB detection+tracking over an ordered frame source -> trajectory CSV.
466
467     Two equivalent ways to supply pixels (output is bit-identical between them):
468
469     * **frames-on-disk** (default): ``frames`` is a directory path string; frames are
470       globbed (``*.png``/``*.jpg``) and ``ImageDataset`` reads each file via
471       ``cv2.imread``.

```

```

471     * **streaming** (fast path): ``frames`` is a pre-built ordered list of (synthetic)
472     frame paths and ``frame_loader(index) -> BGR uint8 ndarray`` supplies the decoded
473     pixels in-memory. We scope a shim over WASB's ``read_image`` (the actual per-frame
474     reader, which returns a PIL RGB image) over the DataLoader pass so every other
line of
475     ``ImageDataset.__getitem__`` runs UNCHANGED ? the tensor the detector sees is
identical
476     to the disk round-trip (``cv2.imwrite`` PNG -> ``Image.open().convert('RGB')`` ==
477     ``cvtColor(bgr, BGR2RGB)``), verified byte-identical). Streaming forces
``num_workers=0``
478     (the in-process shim + buffer can't cross worker processes).
479
480     ``source_key`` overrides the cache-keying identity (defaults to the abspath of the
481     frames dir); the streaming path passes a stable ``video:<abspath>`` key so a resume
482     after reboot still matches.
483     """
484     import time
485
486     import torch
487     from dataloaders import build_img_transforms, build_seq_transforms
488     from dataloaders.dataset_loader import ImageDataset
489     from torch.utils.data import DataLoader
490
491     # NB: this is WASB's OWN local `trackers` module (from its repo, on sys.path inside
the
492     # `wasb` conda env), NOT the Roboflow `trackers` package. The `type: ignore`
suppresses
493     # the mypy false positive where the main env resolves `trackers` to the Roboflow
package
494     # (which has no `build_tracker`). There is NO runtime collision ? this runs in a
separate
495     # subprocess/conda env. Do NOT "clean up" the ignore without removing the Roboflow
dep.
496     from trackers import build_tracker # type: ignore[attr-defined]
497     from utils import Center
498
499     from detectors import build_detector
500
501     streaming = frame_loader is not None
502
503     cfg = _build_cfg(weights, sport, device)
504     frames_in = int(cfg["model"]["frames_in"])
505     input_wh = (int(cfg["model"]["inp_width"]), int(cfg["model"]["inp_height"]))
506     output_wh = (int(cfg["model"]["out_width"]), int(cfg["model"]["out_height"]))
507
508     if streaming:
509         # `frames` is already the ordered list of (synthetic) frame paths.
510         if not isinstance(frames, (list, tuple)):
511             raise SystemExit(
512                 "streaming mode requires `frames` to be a list of frame paths"
513             )
514         frames = list(frames)
515         frames_dir = source_key or "stream"
516     else:
517         frames_dir = frames
518         frames = sorted(
519             glob.glob(osp.join(frames_dir, "*.png"))
520             + glob.glob(osp.join(frames_dir, "*.jpg"))
521         )
522     if limit:
523         frames = frames[:limit]
524     if len(frames) < frames_in:
525         where = frames_dir if not streaming else "video stream"
526         raise SystemExit(f"need >= {frames_in} frames, found {len(frames)} in {where}")
527

```

```

528         # Sliding windows of `frames_in` consecutive frames (the unit of GPU work +
checkpoint).
529         samples = []
530         for window in build_windows(frames, frames_in):
531             annos = [
532                 {"frame_path": p, "center": Center(is_visible=False, x=-1.0, y=-1.0)}
533                 for p in window
534             ]
535             samples.append({"frames": window, "annos": annos})
536         windows_total = len(samples)
537
538         # ---- cache setup / resume decision ----- #
539         if cache_dir is None:
540             cache_dir = default_cache_dir(out_csv)
541             os.makedirs(cache_dir, exist_ok=True)
542             det_jsonl, _ = _cache_paths(cache_dir)
543             cache_key = source_key if source_key is not None else osp.abspath(frames_dir)
544             manifest = None if fresh else read_manifest(cache_dir)
545             resume = manifest_compatible(
546                 manifest or {},
547                 frames_dir=cache_key,
548                 weights=weights,
549                 sport=sport,
550                 frames_in=frames_in,
551                 frames_total=len(frames),
552                 device=device,
553             )
554             if resume:
555                 windows_done, det_by_fid = load_and_clean_cache(cache_dir)
556                 logger.info(
557                     f"resuming from cache: {windows_done}/{windows_total} windows already
detected"
558                 )
559             else:
560                 if manifest is not None:
561                     logger.info("cache incompatible with this run -> starting fresh")
562                     reset_cache(cache_dir)
563                     windows_done, det_by_fid = 0, defaultdict(list)
564
565             base_manifest = {
566                 "version": CACHE_VERSION,
567                 "frames_dir": cache_key,
568                 "weights": weights,
569                 "sport": sport,
570                 "frames_in": frames_in,
571                 "frames_total": len(frames),
572                 "device": device,
573                 "windows_total": windows_total,
574                 "windows_done": windows_done,
575             }
576             write_manifest(cache_dir, base_manifest)
577
578         # ---- detector pass over the un-cached windows ----- #
579         remaining = samples[windows_done:]
580         if remaining:
581             _, transform_test = build_img_transforms(cfg)
582             try:
583                 _, seq_transform_test = build_seq_transforms(cfg)
584             except Exception:
585                 seq_transform_test = None
586             ds = ImageDataset(
587                 cfg,
588                 remaining,
589                 input_wh,
590                 output_wh,

```

```

591         transform=transform_test,
592         seq_transform=seq_transform_test,
593         is_train=False,
594     )
595     # Streaming forces single-process loading: the in-memory frame store + the
596     # cv2.imread shim live in THIS process and can't be pickled to DataLoader
workers.
597     loader_workers = 0 if streaming else num_workers
598     loader = DataLoader(
599         ds, batch_size=batch_size, shuffle=False, num_workers=loader_workers
600     )
601     detector = build_detector(cfg)
602
603     from contextlib import ExitStack
604
605     t0 = time.time()
606     done = windows_done
607     win_idx = windows_done
608     log_every = max(1, log_every_batches)
609     flush_n = max(1, flush_every)
610     pending = []
611     logger.info(
612         f"detecting on {len(remaining)} remaining windows "
613         f"(total {windows_total}, batch={batch_size}, workers={loader_workers}, "
614         f"source={'video-stream' if streaming else 'frames-dir'})..."
615     )
616     det_f = open(det_jsonl, "a")
617     try:
618         with ExitStack() as stack:
619             stack.enter_context(torch.no_grad())
620             # In streaming mode, route ImageDataset's per-frame `cv2.imread(path)`
to the
621             # in-memory decoded frame for that synthetic path; every other line of
622             # __getitem__ runs unchanged so the tensor matches the PNG round-trip
exactly.
623             # The detector's first needed frame is `windows_done` (window i starts
at
624             # frame i), so prime the store there for a resume.
625             if streaming:
626                 stack.enter_context(
627                     _streaming_read_image_patch(
628                         frame_loader,
629                         len(frames),
630                         window=frames_in,
631                         start_index=windows_done,
632                     )
633                 )
634             # Overlap CPU batch assembly with GPU inference (streaming/workers=0
only ?
635             # a multi-worker DataLoader already overlaps via its worker processes).
636             batches = (
637                 _PrefetchIter(loader, depth=prefetch_batches)
638                 if prefetch_batches and loader_workers == 0
639                 else loader
640             )
641             for bi, batch in enumerate(batches):
642                 imgs, _hms, trans = batch[0], batch[1], batch[2]
643                 img_paths = [
644                     list(t) for t in batch[-1]
645                 ] # [frame_in][batch] -> path
646                 batch_results, _ = detector.run_tensor(imgs, trans)
647                 nb = imgs.shape[0]
648                 for ib in range(nb):
649                     frames_payload = []
650                     for ie in sorted(batch_results[ib].keys()):

```

```

651         p = img_paths[ie][ib]
652         fid = _frame_id(p)
653         plain = [_det_plain(d) for d in batch_results[ib][ie]]
654         frames_payload.append([fid, plain])
655         det_by_fid[fid].extend(plain)
656         pending.append({"w": win_idx, "f": frames_payload})
657         win_idx += 1
658     done += nb
659     if len(pending) >= flush_n or done >= windows_total:
660         _append_windows(det_f, pending)
661         pending = []
662         base_manifest["windows_done"] = done
663         write_manifest(cache_dir, base_manifest)
664     if (bi + 1) % log_every == 0 or done >= windows_total:
665         el = time.time() - t0
666         new_done = done - windows_done
667         rate = new_done / el if el > 0 else 0.0
668         eta = (windows_total - done) / rate if rate > 0 else 0.0
669         logger.info(
670             f"{done}/{windows_total} "
671             f"({100 * done // max(1, windows_total)}%) | {rate:.1f}
win/s | "
672             f"elapsed {el:.0f}s | ETA {eta:.0f}s"
673         )
674     if pending:
675         _append_windows(det_f, pending)
676         base_manifest["windows_done"] = done
677         write_manifest(cache_dir, base_manifest)
678     finally:
679         det_f.close()
680         logger.info(f"detection done in {time.time() - t0:.0f}s; running tracker...")
681 else:
682     logger.info("detector cache already complete -> tracker-only re-run")
683
684 # ---- tracker re-run over the full ordered cache (cheap, deterministic) ---- #
685 tracker = build_tracker(cfg)
686 tracker.refresh()
687 rows = []
688 for fid in sorted(det_by_fid.keys()):
689     r = tracker.update(_dets_to_tracker(det_by_fid[fid]))
690     rows.append((fid, 1 if r.get("visi") else 0, r.get("x", -1), r.get("y", -1)))
691
692 _atomic_write_trajectory(out_csv, rows)
693 visible = sum(1 for _, v, *_ in rows if v)
694 logger.info(f"wrote {len(rows)} rows -> {out_csv}")
695 logger.info(f"visible-shuttle frames: {visible}/{len(rows)}")
696 return out_csv
697
698
699 def extract_frames(video_path: str, frames_dir: str) -> str:
700     """Decode a video to PNG frames named by source frame index (00000000.png ...).
701
702     Idempotent: if the dir already holds the expected number of frames (matching the
703     container's reported frame count) we skip re-decoding ? re-extraction after a
704     reboot is wasteful. NOTE: cv2 PNG extraction is slow; for big clips prefer
705     extracting with ffmpeg on the host and passing --frames_dir.
706     """
707     import cv2
708
709     os.makedirs(frames_dir, exist_ok=True)
710     cap = cv2.VideoCapture(video_path)
711     if not cap.isOpened():
712         raise SystemExit(f"cannot open video: {video_path}")
713     expected = int(cap.get(cv2.CAP_PROP_FRAME_COUNT) or 0)
714     existing = len(glob.glob(osp.join(frames_dir, "*.png")))

```

```

715         if expected > 0 and existing >= expected:
716             cap.release()
717             logger.info(f"reusing {existing} already-extracted frames -> {frames_dir}")
718             return frames_dir
719     i = 0
720     while True:
721         ok, frame = cap.read()
722         if not ok:
723             break
724         cv2.imwrite(osp.join(frames_dir, f"{i:08d}.png"), frame)
725         i += 1
726     cap.release()
727     logger.info(f"extracted {i} frames -> {frames_dir}")
728     return frames_dir
729
730
731 def _probe_frame_count(video_path: str) -> int:
732     """Container-reported frame count via cv2 (0 if unknown). Used to size the
733     stream."""
734     import cv2
735     cap = cv2.VideoCapture(video_path)
736     if not cap.isOpened():
737         raise SystemExit(f"cannot open video: {video_path}")
738     n = int(cap.get(cv2.CAP_PROP_FRAME_COUNT) or 0)
739     cap.release()
740     return n
741
742
743 def make_video_frame_loader(video_path: str):
744     """Return ``(frame_loader, count_hint, release)`` for output-preserving streaming.
745
746     ``frame_loader(index) -> BGR uint8 ndarray`` decodes the video forward-only with one
747     long-lived ``cv2.VideoCapture``. It MUST be called with non-decreasing indices (the
748     detector's sliding-window order); it reads sequentially and only uses a frame-peek
749     to
750     skip forward when a *resume* starts past the current position. Each ``cap.read()``
751     returns exactly the BGR uint8 array that ``cv2.imwrite``/``cv2.imread`` of a PNG
752     would
753     yield (PNG is lossless), so the detector sees identical pixels to the disk path.
754
755     Returns the count hint from the container so the caller can build the frame list;
756     ``release`` closes the capture.
757     """
758     import cv2
759     cap = cv2.VideoCapture(video_path)
760     if not cap.isOpened():
761         raise SystemExit(f"cannot open video: {video_path}")
762     count_hint = int(cap.get(cv2.CAP_PROP_FRAME_COUNT) or 0)
763     state = {"pos": 0} # next index cap.read() will return
764
765     def frame_loader(index: int):
766         index = int(index)
767         if index < state["pos"]:
768             raise RuntimeError(
769                 f"streaming decode is forward-only; got index {index} after
770                 {state['pos']}")
771         if index > state["pos"]:
772             # Forward skip only happens when resuming past cached windows. Seek once.
773             cap.set(cv2.CAP_PROP_POS_FRAMES, index)
774             state["pos"] = index
775         ok, frame = cap.read()
776         if not ok:

```

```

776         raise RuntimeError(f"failed to decode frame {index} from {video_path}")
777         state["pos"] += 1
778         return frame
779
780     def release():
781         cap.release()
782
783     return frame_loader, count_hint, release
784
785
786 def run_video_streaming(
787     video_path: str,
788     weights: str,
789     out_csv: str,
790     sport: str = "badminton",
791     limit: int = 0,
792     batch_size: int = 8,
793     log_every_batches: int = 50,
794     cache_dir: "str | None" = None,
795     flush_every: int = 1000,
796     fresh: bool = False,
797     device: str = "cuda",
798     prefetch_batches: int = 0,
799 ) -> str:
800     """Fast path: decode ``video_path`` on the fly and feed frames straight to the
801     detector, WITHOUT extracting every frame to a PNG first (the disk round-trip that
802     dominates a long clip). Output is bit-identical to the frames-on-disk path; see
803     ``run``'s docstring for why.
804     """
805     frame_loader, count_hint, release = make_video_frame_loader(video_path)
806     try:
807         if count_hint <= 0:
808             raise SystemExit(
809                 f"could not determine frame count for {video_path}; cannot stream "
810                 f"(fall back to frame extraction with --frames_dir)"
811             )
812         frame_list = [synthetic_frame_path(i) for i in range(count_hint)]
813         source_key = "video:" + osp.abspath(video_path)
814         return run(
815             frame_list,
816             weights,
817             out_csv,
818             sport,
819             limit,
820             batch_size=batch_size,
821             num_workers=0,
822             log_every_batches=log_every_batches,
823             cache_dir=cache_dir,
824             flush_every=flush_every,
825             fresh=fresh,
826             frame_loader=frame_loader,
827             source_key=source_key,
828             device=device,
829             prefetch_batches=prefetch_batches,
830         )
831     finally:
832         release()
833
834
835 def main():
836     ap = argparse.ArgumentParser()
837     ap.add_argument("--frames_dir", help="folder of pre-extracted frames")
838     ap.add_argument(
839         "--video", help="video file; frames are extracted internally (in WSL)"
840     )

```

```

841 ap.add_argument(
842     "--frames_out_dir", help="where to extract frames when --video is used"
843 )
844 ap.add_argument(
845     "--stream-video",
846     action="store_true",
847     help="FAST PATH: decode --video on the fly and feed frames straight to "
848     "the detector (no PNG extraction). Output-equivalent to the disk path.",
849 )
850 ap.add_argument("--weights", required=True)
851 ap.add_argument("--out", required=True)
852 ap.add_argument("--sport", default="badminton")
853 ap.add_argument(
854     "--device",
855     default="cuda",
856     choices=["cuda", "mps", "cpu"],
857     help="torch device for the detector (default cuda = the WSL parity path)",
858 )
859 ap.add_argument(
860     "--limit", type=int, default=0, help="cap #frames (for quick tests)"
861 )
862 ap.add_argument("--batch-size", type=int, default=8)
863 ap.add_argument("--num-workers", type=int, default=4)
864 ap.add_argument(
865     "--prefetch-batches",
866     type=int,
867     default=0,
868     help="streaming mode: overlap CPU batch assembly with GPU inference via a "
869     "bounded producer thread of this queue depth (0 = off, the parity default)",
870 )
871 ap.add_argument(
872     "--log-every-batches", type=int, default=50, help="progress log cadence"
873 )
874 ap.add_argument(
875     "--cache-dir",
876     default=None,
877     help="resumable detector cache dir (default: <out>_wasbcache next to --out)",
878 )
879 ap.add_argument(
880     "--flush-every",
881     type=int,
882     default=1000,
883     help="fsync the detector cache every N windows (bounds reboot loss)",
884 )
885 ap.add_argument(
886     "--fresh", action="store_true", help="ignore any existing cache and start over"
887 )
888 args = ap.parse_args()
889
890 if args.stream_video:
891     if not args.video:
892         raise SystemExit("--stream-video requires --video")
893     run_video_streaming(
894         args.video,
895         args.weights,
896         args.out,
897         args.sport,
898         args.limit,
899         batch_size=args.batch_size,
900         log_every_batches=args.log_every_batches,
901         cache_dir=args.cache_dir,
902         flush_every=args.flush_every,
903         fresh=args.fresh,
904         device=args.device,
905         prefetch_batches=args.prefetch_batches,

```



```

906         )
907         return
908
909     frames_dir = args.frames_dir
910     if args.video:
911         frames_dir = extract_frames(
912             args.video, args.frames_out_dir or (args.video + "_frames")
913         )
914     if not frames_dir:
915         raise SystemExit("provide --frames_dir or --video")
916     run(
917         frames_dir,
918         args.weights,
919         args.out,
920         args.sport,
921         args.limit,
922         batch_size=args.batch_size,
923         num_workers=args.num_workers,
924         log_every_batches=args.log_every_batches,
925         cache_dir=args.cache_dir,
926         flush_every=args.flush_every,
927         fresh=args.fresh,
928         device=args.device,
929     )
930
931
932 if __name__ == "__main__":
933     main()
934

```

4. user (2026-07-03T19:48:27.722Z)

```

1     """Linux-native WASB shuttle-trajectory runner (no WSL).
2
3     The de-risk goal the ``DetectorRunner`` ABC was built for (see ``base.py`` ?
4     "Optional Linux-native path"; ``docs/DECOUPLED_COMPUTE/04-BUILD-AND-TEST-PLAN.md``
5     M3.5).
6     It is the sibling of :class:`~backend.pipeline.detectors.wasb_runner.WasbRunner` for a
7     Linux GPU box (GCP / a SaaS worker): it runs the same verified ``wasb_infer.py``
8     core
9     that the WSL adapter runs, so the trajectory CSV is identical (to a sub-pixel float
10    tolerance;
11    see the acceptance gate below) ? but WITHOUT the WSL transport:
12
13    * NO ``WslCommandMixin`` (no ``wsl -d <distro> bash -c``), NO ``conda activate`` (the
14    ``wasb`` env's python is invoked by path), NO ``/mnt/c`` path translation
15    (``to_wsl_mnt_path``). Paths are native; the video is already local on the box.
16    * ``device`` is ``cuda`` (default ? the WSL parity path), ``mps``, ``cpu``, or
17    ``auto``
18    (M4: cuda-if-available-else-cpu ? the CPU fallback for a no-GPU box; explicit
19    ``cuda`` stays
20    strict); threaded into ``wasb_infer.py --device`` (and so the Hydra
21    ``runner.device`` override).
22    * weights / repo / python / scratch come from the environment first
23    (``WASB_WEIGHTS`` / ``WASB_REPO_DIR`` / ``WASB_PYTHON`` / ``WASB_SCRATCH``), then
24    ``indexing.wasb.*`` config, then a default ? so a box needs no config edits.
25
26    It emits the identical ``Frame,Visibility,X,Y`` CSV (same ``{stem}_{vid12}_wasb.csv``
27    name as the WSL runner) and reuses the model-agnostic ``parse_trajectory_csv`` /
28    ``trajectory_to_action_windows`` from ``tracknet_runner``, so the trajectory hybrids and
29    the windowing brain are unchanged. Selected via the existing ``detector_impl`` seam
30    (``trajectory_hybrid._resolve_runner`` ? ``_build_native_runner``);
31    ``detector_impl="native"``.
32
33    Acceptance gate (parity, NOT strict byte-equality): the SAME clip through the WSL

```

```

runner
27     (Windows) and this runner (Linux, ``device=cuda``) yields the same trajectory CSV + the
same
28     candidate windows **within a small float tolerance (~1e-3 px)**. cuDNN is non-
deterministic
29     ACROSS GPUs, so a sub-pixel diff is expected on different hardware ? but it is byte-
exact
30     run-to-run on a FIXED GPU. ? Validated 2026-06-20 on an RTX 2070 (via WSL): native
streaming
31     vs the cached WSL-disk CSV matched **1194/1195** frames (the lone diff a ~5e-5 px edge
32     artifact), and two native runs were byte-identical (deterministic). This is a hardware
check
33     (a real GPU + the ``wasb`` env), so the offline, subprocess-mocked suite asserts the
contract,
34     not the numbers.
35     ""
36
37     import logging
38     import os
39     import shutil
40     import subprocess
41     from dataclasses import dataclass
42     from typing import List, Optional, Tuple
43
44     from backend.config.models import IndexingConfig
45     from backend.pipeline.detectors.base import DetectorRunner
46     from backend.pipeline.detectors.device import AUTO, VALID_DEVICES, DeviceContext
47
48     # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical to
WSL/TrackNet.
49     from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
50
51     logger = logging.getLogger(__name__)
52
53     # Path to the inference core we copy into the WASB repo's src/ (so it can import WASB's
54     # detectors/trackers/dataloaders + find configs/), exactly as the WSL adapter does.
55     _INFER_PY = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")
56
57     # Single torch/cuda probe code, used by BOTH healthcheck and the lazy device-resolution
probe
58     # so the "cuda True/False" string they parse can never drift apart.
59     _CUDA_PROBE_CODE = (
60         "import torch; print('torch', torch.__version__, 'cuda', torch.cuda.is_available())"
61     )
62
63
64     @dataclass
65     class NativeWasbConfig:
66         """Resolved settings for a Linux-native WASB run.
67
68         Built via :meth:`from_indexing_cfg` with **environment overrides** taking precedence
69         over ``indexing.wasb.*`` config (the box sets env; no config edits needed).
``device``
70         defaults to ``cuda`` (the WSL parity path); ``stream_video`` defaults to True (the
perf win).
71         ""
72
73         repo_dir: str = "~/models/WASB-SBDT" # native clone of nttcom/WASB-SBDT
74         python_bin: str = (
75             "python" # the `wasb` conda env python (invoked by path, no activate)
76         )
77         weights_path: str = "" # REQUIRED (env WASB_WEIGHTS or indexing.wasb.weights_path)
78         sport: str = "badminton"
79         device: str = "cuda" # auto | cuda | mps | cpu (auto = cuda-if-available-else-cpu)
80         scratch_dir: str = "" # frame-extraction scratch (env); "" = next to the output dir

```

```

81     timeout_sec: int = 1800 # 30 min; a hung GPU call is killed (0 = no timeout)
82     keep_frames: bool = (
83         False # retain the decoded frame cache after success (debug only)
84     )
85     # NATIVE DEFAULT = STREAMING: skip the cv2 PNG extraction that dominates a long
clip's
86     # wall-clock (measured ~50s on a 20s clip) ? bit-identical to disk (verified on a
real GPU
87     # 2026-06-20). The GPU box reads from a local NVMe so streaming is the right default
here.
88     stream_video: bool = True
89     # GPU-feed tuning (indexing.wasb.batch_size / prefetch_batches). 0 = omit the flag ?
90     # wasb_infer.py's parity defaults (batch 8, no prefetch). The cloud-serving preset
opts in.
91     batch_size: int = 0
92     prefetch_batches: int = 0
93
94     @classmethod
95     def from_indexing_cfg(cls, idx_cfg: "IndexingConfig") -> "NativeWasbConfig":
96         if isinstance(idx_cfg, dict):
97             idx_cfg = IndexingConfig(**idx_cfg)
98         w = idx_cfg.wasb
99         env = os.environ.get
100         # stream_video is tri-state in config: None/absent => native default (stream);
an
101         # explicit True/False still wins (e.g. False for a container cv2 can't seek).
102         sv = w.stream_video
103         return cls(
104             repo_dir=env("WASB_REPO_DIR") or w.repo_dir,
105             python_bin=env("WASB_PYTHON") or w.python_bin,
106             weights_path=env("WASB_WEIGHTS") or w.weights_path or "",
107             sport=w.sport,
108             device=env("WASB_DEVICE") or w.device,
109             scratch_dir=env("WASB_SCRATCH")
110             or env("RALLY_SCRATCH")
111             or env("TMPDIR")
112             or "",
113             timeout_sec=w.timeout_sec,
114             keep_frames=w.keep_frames,
115             stream_video=(True if sv is None else bool(sv)),
116             batch_size=int(w.batch_size or 0),
117             prefetch_batches=int(w.prefetch_batches or 0),
118         )
119
120
121     class NativeWasbRunner(DetectorRunner):
122         """WASB runner that executes natively on a Linux GPU box ? no WSL. See module
docstring."""
123
124         def __init__(self, cfg: NativeWasbConfig):
125             self.cfg = cfg
126             # Cached CUDA-availability probe (used to resolve device="auto"). Set by
healthcheck;
127             # lazily probed by run_predict if healthcheck was skipped. None = not yet
probed.
128             self._cuda_available: Optional[bool] = None
129
130         # --- low-level native exec (the single place tests patch)
-----
131         def _run(
132             self, argv: List[str], cwd: Optional[str] = None
133         ) -> subprocess.CompletedProcess:
134             logger.debug("native exec: %s (cwd=%s)", " ".join(argv), cwd)
135             return subprocess.run(
136                 argv,

```

```

137         cwd=cwd,
138         capture_output=True,
139         text=True,
140         timeout=(self.cfg.timeout_sec or None),
141     )
142
143     # --- device resolution (M4 DeviceContext: device="auto" ? cuda-if-available-else-
cpu) -----
144     def _probe_cuda(self) -> bool:
145         """Probe ``torch.cuda.is_available()`` in the wasb env's python (a subprocess).
Any failure
146         (missing python / timeout / import error) is treated as 'no CUDA'."""
147         try:
148             res = self._run([self.cfg.python_bin, "-c", _CUDA_PROBE_CODE])
149         except (FileNotFoundError, subprocess.TimeoutExpired):
150             return False
151         return res.returncode == 0 and "cuda True" in (res.stdout or "")
152
153     def _cuda_is_available(self) -> bool:
154         """CUDA availability, cached across healthcheck + run_predict so we probe at
most once."""
155         if self._cuda_available is None:
156             self._cuda_available = self._probe_cuda()
157         return self._cuda_available
158
159     def _effective_device(self) -> str:
160         """The torch device to actually pass to ``wasb_infer.py``. Only ``auto`` needs
the CUDA
161         probe; an explicit cuda/mps/cpu resolves to itself (so ``device='cuda'`` is
unchanged)."""
162         if self.cfg.device == AUTO:
163             return DeviceContext(AUTO, self._cuda_is_available()).effective
164         return self.cfg.device
165
166     def _repo_src(self) -> str:
167         return os.path.join(os.path.expanduser(self.cfg.repo_dir), "src")
168
169     def _copy_infer_script(self) -> bool:
170         """Copy wasb_infer.py into the WASB repo src/ so it imports WASB modules + finds
configs/."""
171         repo_src = self._repo_src()
172         try:
173             os.makedirs(repo_src, exist_ok=True)
174             shutil.copy(_INFER_PY, os.path.join(repo_src, "wasb_infer.py"))
175             return True
176         except OSError as e:
177             logger.error("Failed to copy wasb_infer.py into %s: %s", repo_src, e)
178         return False
179
180     def _rm_rf(self, path: Optional[str]) -> None:
181         """Best-effort recursive delete of a native path (post-success cleanup; never
raises)."""
182         if path:
183             shutil.rmtree(path, ignore_errors=True)
184
185     def healthcheck(self) -> Tuple[bool, str]:
186         """Verify weights + repo present and the `wasb` python imports torch (and sees a
GPU on cuda)."""
187         if not self.cfg.weights_path:
188             return False, (
189                 "WASB weights path is empty ? set WASB_WEIGHTS (or
indexing.wasb.weights_path)."
190             )
191         weights = os.path.expanduser(self.cfg.weights_path)
192         if not os.path.exists(weights):

```

```

193         return False, f"WASB weights not found at {weights}"
194     repo_src = self._repo_src()
195     if not os.path.isdir(repo_src):
196         return False, (
197             f"WASB-SBDT repo src/ not found at {repo_src} ? clone nttcom/WASB-SBDT "
198             f"(set WASB_REPO_DIR / indexing.wasb.repo_dir).")
199     )
200     # Fail fast on a typo'd device (the device-policy seam) ? a clear error here
beats a cryptic
201     # argparse rc=2 from wasb_infer.py at run time.
202     if self.cfg.device not in VALID_DEVICES:
203         return False, (
204             f"unknown device {self.cfg.device!r} ? valid: {'',
205             ' '.join(VALID_DEVICES)}").
206     )
207     try:
208         res = self._run([self.cfg.python_bin, "-c", _CUDA_PROBE_CODE])
209     except FileNotFoundError:
210         return (
211             False,
212             f"python binary not found: {self.cfg.python_bin!r} (set WASB_PYTHON).",
213         )
214     except subprocess.TimeoutExpired:
215         return False, "native torch healthcheck timed out."
216     out = (res.stdout or "").strip()
217     if res.returncode != 0:
218         return (
219             False,
220             f"`{self.cfg.python_bin}` torch import failed: {(res.stderr or
221             out).strip()}",
222         )
223     # M4: resolve the device. Cache the probe so run_predict reuses it (no second
224     subprocess).
225     self._cuda_available = "cuda True" in out
226     ctx = DeviceContext(self.cfg.device, self._cuda_available)
227     if ctx.cuda_required_but_missing:
228         # Unchanged pre-M4 behaviour for an EXPLICIT cuda request: hard-fail, never
a silent
229         # slow-CPU downgrade. Set indexing.wasb.device=auto for a CPU fallback.
230         return False, f"device=cuda but torch.cuda.is_available() is False ({out})"
231     if ctx.fell_back:
232         logger.warning(
233             "device=auto: no CUDA GPU detected (%s) ? FALLING BACK TO CPU; WASB "
234             "inference will be slow. Set indexing.wasb.device=cuda to require a
GPU.",
235             out,
236         )
237     return True, out
238
239     def run_predict(
240         self, video_win_path: str, output_win_dir: str, video_id: Optional[str] = None
241     ) -> Optional[str]:
242         """Run WASB natively on a video. Returns the path to the trajectory CSV, or
None.
243
244         Output artifacts are namespaced by ``video_id`` (the file MD5) exactly as the
WSL
245         runner does, and the CSV is named ``{stem}_{vid12}_wasb.csv`` ? byte-for-byte
the
246         same downstream contract, so the parity comparison is apples-to-apples.
247         """
248         output_dir = os.path.abspath(output_win_dir)
249         os.makedirs(output_dir, exist_ok=True)
250         # Normalize for a cross-platform basename (tests pass Windows paths on Linux).
251         norm = str(video_win_path).replace("\\", "/")

```

```

249         stem, _ext = os.path.splitext(os.path.basename(norm))
250         # Content-derived key: same filename + different bytes -> different key (no
cache reuse).
251         key = f"{stem}__{video_id[:12]}" if video_id else (stem or "wasb")
252         out_csv = os.path.join(output_dir, f"{key}_wasb.csv")
253
254         if not self._copy_infer_script():
255             logger.error("Could not stage wasb_infer.py into the WASB repo src/.")
256             return None
257
258         weights = os.path.expanduser(self.cfg.weights_path)
259         argv = [
260             self.cfg.python_bin,
261             "wasb_infer.py",
262             "--video",
263             os.path.abspath(str(video_win_path)),
264         ]
265         frames_dir: Optional[str] = None
266         if self.cfg.stream_video:
267             # Fast path: decode in-process, no PNG extraction. Output is bit-identical
to disk.
268             argv.append("--stream-video")
269         else:
270             scratch = self.cfg.scratch_dir or output_dir
271             frames_dir = os.path.join(scratch, f"{key}_frames")
272             argv += ["--frames_out_dir", frames_dir]
273         # M4: resolve device="auto" to the effective device (cuda-if-available-else-
cpu); an
274         # explicit cuda/mps/cpu passes through unchanged. Never sends "auto" to
wasb_infer.py.
275         device = self._effective_device()
276         argv += [
277             "--weights",
278             weights,
279             "--sport",
280             self.cfg.sport,
281             "--device",
282             device,
283             "--out",
284             out_csv,
285         ]
286         # GPU-feed tuning: pass only explicit non-zero values so the default argv (and
its
287         # byte-parity guarantee) is unchanged when the knobs are unset.
288         if self.cfg.batch_size:
289             argv += ["--batch-size", str(self.cfg.batch_size)]
290         if self.cfg.prefetch_batches:
291             argv += ["--prefetch-batches", str(self.cfg.prefetch_batches)]
292
293         logger.info(
294             "Running native WASB inference (device=%s%s) on %s ...",
295             device,
296             " [auto?fallback]" if self.cfg.device == AUTO and device != "cuda" else "",
297             os.path.basename(norm),
298         )
299         try:
300             res = self._run(argv, cwd=self._repo_src())
301         except FileNotFoundError:
302             logger.error(
303                 "python binary not found: %r (set WASB_PYTHON).", self.cfg.python_bin
304             )
305             return None
306         except subprocess.TimeoutExpired:
307             logger.error(
308                 "WASB native inference timed out after %ss.", self.cfg.timeout_sec

```

```

309         )
310         return None
311     if res.returncode != 0:
312         logger.error(
313             "WASB native inference failed:\n%s",
314             (res.stderr or res.stdout or "")[-2000:],
315         )
316         return None
317     if not os.path.exists(out_csv):
318         logger.error(
319             "WASB native finished but expected CSV not found at %s", out_csv
320         )
321         return None
322     logger.info("WASB trajectory CSV: %s", out_csv)
323
324     # Success ? the CSV is the durable output; the decoded frame cache is a pure
intermediate.
325     # On any earlier failure we return above WITHOUT cleaning, so the cache survives
a resume.
326     # (Streaming never materializes frames, so there is nothing to drop on that
path.)
327     if not self.cfg.keep_frames and frames_dir:
328         logger.info(
329             "Cache hygiene: removing frame cache for %s "
330             "(set indexing.wasb.keep_frames=true to retain).",
331             key,
332         )
333         self._rm_rf(frames_dir)
334     return out_csv
335
336     # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
337     parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
338     trajectory_to_action_windows = staticmethod(
339         TrackNetRunner.trajectory_to_action_windows
340     )
341

```

5. assistant (2026-07-03T19:48:30.324Z)

Now let me examine the tracker re-run loop closely. Let me look at the key sections and understand the data structures involved. Let me check the config models and the tracker's typical implementation pattern.

6. user (2026-07-03T19:48:33.934Z)

No files found

7. user (2026-07-03T19:48:34.046Z)

```

1     """Pydantic models for the application configuration.
2
3     This module defines the canonical, validated, typed configuration for the
4     Badminton Highlight Indexer. Built-in defaults live here (single source of truth).
5
6     Unknown / typo'd config keys are rejected at load time via extra="forbid".
7     """
8
9     from __future__ import annotations
10
11     from typing import Any, List, Literal, Optional
12
13     from pydantic import BaseModel, Field, field_validator, model_validator
14

```

```

15     Sport = Literal["badminton", "tennis", "pickleball", "table_tennis", "cricket"]
16
17
18     class ValidationRelevanceConfig(BaseModel):
19         court_green_lower: list[int] = Field(default=[35, 40, 40])
20         court_green_upper: list[int] = Field(default=[85, 255, 255])
21         court_pixels_min_ratio: float = Field(default=0.05, ge=0.0, le=1.0)
22
23
24     class ValidationConfig(BaseModel):
25         min_resolution_width: int = 1280
26         min_resolution_height: int = 720
27         min_fps: float = 29.9
28         min_blur_threshold: float = 20.0
29         max_scene_cuts_per_minute: float = 0.5
30         pts_max_gap_seconds: float = Field(
31             default=10.0,
32             ge=0.0,
33             description=(
34                 "Maximum allowed keyframe gap (seconds) before flagging as spliced. "
35                 "GoPro H.264/H.265 recordings typically have GOP intervals of 1?4s; "
36                 "a gap >10s suggests edited/missing content. Mobile phones, screen "
37                 "recordings, and other encoders may use different GOP sizes ? tune "
38                 "this threshold for your recording device."
39             ),
40         )
41         relevance: ValidationRelevanceConfig = Field(
42             default_factory=ValidationRelevanceConfig
43         )
44
45
46     class TrackNetConfig(BaseModel):
47         wsl_distro: str = "Ubuntu"
48         repo_dir: str = "~/models/TrackNetV4"
49         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
50         conda_env: str = "TrackNetV4"
51         weights_path: str = ""
52         queue_length: int = 5
53         stage_in_wsl: bool = True
54         wsl_stage_dir: str = "~/clips"
55         timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
56         velocity_threshold: float = 0.02
57         # Cache hygiene: after a SUCCESSFUL run the staged video copy (and, for WASB, the
58         # decoded frame cache) are deleted automatically ? they are pure intermediates,
59         # regenerable from the source, and the dominant disk consumer (~GBs/video). Set
60         # true only for debugging. On FAILURE they are always kept so a re-run can resume.
61         keep_frames: bool = False
62         # Absolute FLOOR (GB) on WSL free space (at wsl_stage_dir) before a run, on top of
63         # ~15x-source headroom the frame-stage guard always enforces (D11 ? it now BLOCKS,
64         # 0 = no extra floor (the headroom guard still applies).
65         min_free_gb: float = 0.0
66
67
68     class WasbIndexConfig(BaseModel):
69         """Typed config for the live WASB shuttle substrate (indexing.wasb).
70
71         Mirrors backend/pipeline/detectors/wasb_runner.py::WasbConfig.from_indexing_cfg so
72         these knobs actually take effect ? previously this whole block was silently dropped
73         because nested config sections defaulted to extra='ignore'.
74         """
75
76         wsl_distro: str = "Ubuntu"
77         repo_dir: str = "~/models/WASB-SBDT"

```



```

78     conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
79     conda_env: str = "wasb"
80     weights_path: str = (
81         "~/models/WASB-SBDT/pretrained_weights/wasb_badminton_best.pth.tar"
82     )
83     sport: str = "badminton"
84     wsl_stage_dir: str = "~/clips"
85     timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
86     velocity_threshold: float = 0.02
87     # --- Linux-native runner only (detector_impl='native'); ignored by the WSL runner.
88     # Env vars override these (WASB_DEVICE / WASB_PYTHON) for a config-free GPU box.
89     # torch device for the native runner. "auto" = cuda-if-available-else-cpu (the M4
CPU fallback
90     # for a no-GPU box); "cuda" stays strict (hard-fails without a GPU ? no silent slow-
CPU downgrade).
91     device: str = "cuda" # auto | cuda | mps | cpu
92     python_bin: str = (
93         "python" # the `wasb` conda env python (invoked by path, no `conda activate`)
94     )
95     # Cache hygiene: after a SUCCESSFUL run the staged video copy + decoded frame cache
96     # (the ~GBs/video disk hog) are deleted automatically ? pure intermediates,
regenerable
97     # from the source. Set true only for debugging. On FAILURE they are kept so a re-run
resumes.
98     keep_frames: bool = False
99     # Absolute FLOOR (GB) on WSL free space (at wsl_stage_dir) before a run, on top of
the
100    # ~15x-source headroom the frame-stage guard always enforces (D11 ? it now BLOCKS,
not warns).
101    # 0 = no extra floor (the headroom guard still applies).
102    min_free_gb: float = 0.0
103    # FAST PATH ? decode the video on the fly and feed frames straight to the detector,
104    # skipping the slow per-frame PNG extraction that dominates a long clip (~46k PNGs
for a
105    # 12-min 1080p60 video, which overran the 30-min timeout). Output is bit-identical
to the
106    # PNG path (PNG is lossless) ? VERIFIED on a real GPU 2026-06-20 (native stream ==
WSL disk,
107    # 1194/1195 frames byte-identical; native run-to-run byte-identical /
deterministic).
108    # Tri-state:
109    # None = per-runner default ? the WSL runner keeps DISK extraction
(unchanged Windows
110    # behaviour); the Linux-native runner (GPU box) defaults to STREAMING
(the perf win).
111    # True/False = force on/off for BOTH runners (e.g. set False for a container cv2
can't seek).
112    stream_video: Optional[bool] = None
113    # GPU-feed tuning for the native runner (measured on the first real L4 serving run,
114    # 2026-07-03: the default batch 8 with single-process loading left the GPU <1%
utilized ?
115    # the pipeline was CPU-feed-bound). None = wasb_infer.py's parity defaults (batch 8,
no
116    # prefetch ? byte-identical to the historical path). The cloud-serving preset opts
into
117    # the measured fast values; local/WSL behaviour is unchanged unless set explicitly.
118    batch_size: Optional[int] = None
119    # Depth of the bounded producer-thread queue that overlaps CPU-side batch assembly
120    # (decode + PIL + ToTensor) with GPU inference in streaming mode. None/0 = off.
121    prefetch_batches: Optional[int] = None
122
123
124    class FusionConfig(BaseModel):
125        """Config for the multi-signal `fusion_hybrid` segmenter (MULTI_SIGNAL_FUSION_PLAN

```

```

P2).
126
127 Every default reproduces control behaviour: the fusion segmenter persists the
128 shuttle net-crossing count (Gate A signal) and ? only when ``cue_flags['person']``
is
129 true ? the person-cue features, but it does NOT gate the served handoff unless a
130 threshold is raised (``min_crossings``/``min_players`` default 0 = OFF). The court
mask
131 is optional: ``court_polygon=None`` ? Gate B counts every detection (player-count-
only);
132 supplied ? off-court persons (spectators / adjacent courts) are excluded.
133 ""
134
135 # Which trajectory substrate seeds the windows (shuttle stays primary).
136 substrate: Literal["wasb", "tracknet"] = "wasb"
137 # Windowing velocity threshold for the fusion family (the engine reads
138 # indexing.<family>.velocity_threshold). Default matches the substrates' 0.02; set
139 # equal to indexing.wasb.velocity_threshold to A/B fusion against a tuned wasb run.
140 velocity_threshold: float = 0.02
141 # Per-cue enable flags, e.g. {"person": true}. Person cue is OFF unless explicitly
142 # enabled ? so the default path never imports torch / touches the GPU.
143 cue_flags: dict[str, bool] = Field(default_factory=dict)
144 # Served Gate A: drop windows whose shuttle net-crossings < this from the AI
handoff.
145 # 0 = OFF (no gating; persisted candidates are always the full superset for offline
re-scoring).
146 min_crossings: int = Field(default=0, ge=0)
147 # Served Gate B: when the person cue is on, drop windows with median in-court
players
148 # < this. 0 = OFF.
149 min_players: int = Field(default=0, ge=0)
150 # Optional court mask as normalised 0-1 [[x, y], ...] polygon. None = no mask.
151 court_polygon: Optional[list[list[float]]] = None
152 # Net line for the person two-sided-motion split (person features only ? the shuttle
153 # net-crossing gate keeps rally_gate's camera-agnostic auto-estimate).
154 net_axis: Literal["x", "y"] = "y"
155 net_position: float = Field(default=0.5, ge=0.0, le=1.0)
156
157
158 class IndexingConfig(BaseModel):
159     default_segmenter: str = "yolo_hybrid"
160     use_gpu: bool = True
161     # Detector backend for the trajectory hybrids (wasb/tracknet). "auto" = the real
162     # WSL runner (default, production); "stub" = a deterministic CPU mock (no GPU/WSL)
163     # so the whole pipeline is regression-testable on a CPU box / CI; "native" = the
164     # Linux-native WASB runner (no WSL/conda-activate; a GPU box ? M3.5). The
165     # RALLY_DETECTOR_IMPL env var overrides this. See
pipeline/detectors/{stub,native_wasb}_runner.py.
166     detector_impl: str = "auto"
167     motion_threshold: float = 0.08
168     min_segment_duration: float = 3.0
169     dead_time_merge_gap: float = 2.0
170     chunk_duration: float = 90.0
171     chunk_overlap: float = 10.0
172     detector_model: str = "fasterrcnn_resnet50_fpn"
173     detector_score_threshold: float = 0.5
174     yolo_velocity_threshold: float = 0.15
175     yolo_frame_skip: int = 3
176     yolo_tracker: str = "bytetrack.yaml"
177     yolo_prefetch_workers: int = 2
178     min_action_window_duration: float = 2.0
179     # BOUNDED WINDOWING (over-merge fix W1, RALLY_QUALITY_RESEARCH.md §6). 0 = OFF; > 0
=a window
180     # longer than this many seconds is recursively split at its largest internal
inactivity gap (?)

```

```

181      # ``window_min_split_gap`` s ? the most likely rally boundary), so one window can't
swallow
182      # several rallies + the dead time between them. Never merges, only cuts (recall
can't drop).
183      # *** R1 MILESTONE DEFAULT-ON (cap=6): the smallest safe local-windowing flip.
Measured n=13
184      # golden, R2 tolF1: baseline?milestone (cap=6 + R4 below) ?+0.222, sign 12/0,
p=0.0005; cap=6
185      # beats cap=12 ?+0.110, sign 12/0; net over-seg guard PASS (mean merge_rate
0.651?0.161). ***
186      max_window_duration: float = 6.0
187      window_min_split_gap: float = (
188          0.5 # min internal gap (s) that qualifies as a split point
189      )
190      # HARD-CAP fallback for W1: when True, an over-long window with NO internal
inactivity gap
191      # (a continuously-active trajectory ? e.g. the real-footage mahadevpura-1 174s blob)
is
192      # chopped at its midpoint until ? max_window_duration, making the cap a true upper
bound.
193      # Only cuts (recall can't drop). OFF by default until measured/promoted.
194      window_hard_cap: bool = False
195      # R4 rally-END inactivity trim (s, 0 = OFF). After a rally the shuttle keeps moving
slowly
196      # (picked up / walked) above velocity_thresh, so the window over-extends its END
(the measured
197      # [?end] gap vs golden). Trim the post-rally tail back to the last frame whose
trailing
198      # window stays dense with FAST motion (velocity > inactivity_rally_thresh). Only
cuts.
199      # *** R1 MILESTONE DEFAULT-ON (1.5/0.20/0.30): on top of cap=6, R4 adds a small but
significant
200      # end-precision gain ? ?tolF1 +0.040, sign 9/1/3, p=0.022, [?end] 4.96?3.31s (n=13).
The W1 cap
201      # is the dominant lever; R4 is the marginal increment. See QUALITY_ITERATIONS.md
"R1". ***
202      window_inactivity_end: float = 1.5
203      inactivity_rally_thresh: float = 0.20
204      inactivity_min_density: float = 0.30
205      # R4 density DENOMINATOR fix (code-review finding B). The R4 trim's fast-fraction
denominator is
206      # the COUNT of active samples in the trailing window, which sparse tracking inflates
(one fast
207      # sample after a gap ? 100% dense ? tail held open). When True, the denominator is
the window's
208      # TEMPORAL length instead. No-op under dense tracking (sample-count == temporal-
span), so it only
209      # corrects the sparse-track case. *** P2b PROMOTED DEFAULT-ON (2026-06-29): golden-
set A/B at the
210      # SERVED operating point (n=15, 595 rallies) ? tolF1 0.353?0.364 (?tolF1 +0.010,
sign 10W/4L, NO
211      # regression), [?end] 2.97?2.72s, seg_count_ratio ?0.103 (p=0.006), split_rate
?0.010 (p=0.016),
212      # R1 over-seg guard PASS; served_regresses(eps=0.0)=False ?
promote_if_served_safe(structural)=True.
213      # See docs/RALLY_DETECTION_FIXES_PLAN.md §7/§8. ***
214      inactivity_temporal_density: bool = True
215      # R5 blob-fallback (default OFF). LAST-RESORT splitter for the continuously-active
blob: after
216      # the W1 gap-split AND the R4 inactivity trim have each had their chance, a group
STILL longer
217      # than window_blob_factor × max_window_duration (no internal gap, no rally-end
signal ? e.g.
218      # mahadevpura-1's ~65s residual ? 11× a 6s cap) is midpoint-chopped to ? the cap as
a FORCED cut

```

```

219         # (logged + flagged on the window) so the degenerate single-window outcome is
avoided and the
220         # clip is surfaced as needing rally-STATE cues, not a bigger cap. Differs from
window_hard_cap
221         # (which chops BEFORE R4 and over-segments normal clips): blob-fallback runs AFTER
R4 and only
222         # force-cuts the genuine blob (the factor gate). Only cuts (recall can't drop).
223         window_blob_fallback: bool = False
224         # Magnitude gate that keeps R5 surgical: only a group longer than this ×
max_window_duration is a
225         # blob (a normal long rally ~1-2× the cap is left alone; the mahadevpura-1 residual
is ~11×).
226         # Default 4.0 is do-no-harm on the golden corpus (rescues only the continuously-
active
227         # clips, every other clip a no-op within noise) ? see docs/QUALITY_ITERATIONS.md
"R5".
228         window_blob_factor: float = 4.0
229         log_yolo_candidates: bool = True
230         store_yolo_candidates: bool = True
231         ai_handoff_padding: float = 2.0
232         ai_max_workers: int = 5
233         # Width (px) the gemini segmenter re-encodes chunks to before upload. Stream-copied
234         # chunks of a high-bitrate (4K) source blow past the provider upload cap
235         # (backend/providers/gemini.py::MAX_UPLOAD_MB) and are refused ? observed live as
236         # "0 segments / success". Re-encoding also makes chunk boundaries frame-accurate
237         # (stream copy cuts at the nearest keyframe). 0 = legacy stream-copy at source res.
238         # Matches gemini_refine's --clip-scale-width default.
239         ai_chunk_scale_width: int = 1280
240         # FULLY-LOCAL mode (default OFF): when true, the trajectory-hybrid segmenters
241         # (wasb_hybrid / tracknet_hybrid / fusion_hybrid) skip the Phase-3 AI handoff
entirely and
242         # emit their local candidate windows AS the rallies ? no provider, no API key,
nothing
243         # uploaded. Lets the local detector run standalone (e.g. to benchmark it against the
Gemini
244         # path, or to operate with zero cloud dependency). The pure `gemini` segmenter
ignores this.
245         skip_ai_handoff: bool = False
246         # Optional debug knob: trim every video to the first N seconds before processing.
247         debug_duration: Optional[float] = None
248
249         tracknet: TrackNetConfig = Field(default_factory=TrackNetConfig)
250         wasb: WasbIndexConfig = Field(default_factory=WasbIndexConfig)
251         fusion: FusionConfig = Field(default_factory=FusionConfig)
252
253         @field_validator("default_segmenter")
254         @classmethod
255         def segmenter_must_be_registered(cls, v: str) -> str:
256             # Lazy import: the segmenter registry imports modules that may import
257             # config, so we resolve it at validation time rather than at module load.
258             from backend.pipeline.segmenters.base import segmenter_registry
259
260             if segmenter_registry.get(v) is None:
261                 available = sorted(segmenter_registry.list_available().keys())
262                 raise ValueError(
263                     f"Segmenter '{v}' is not registered. Available: {'', ' '.join(available)}"
264                 )
265             return v
266
267         @model_validator(mode="after")
268         def _chunk_step_must_be_positive(self) -> "IndexingConfig":
269             # The chunked segmenters advance by (chunk_duration - chunk_overlap); a non-
positive
270             # step makes `while t < duration: t += step` loop forever, growing memory before
any

```

```

271         # GPU work. Reject the bad config at load time instead.
272         if self.chunk_overlap >= self.chunk_duration:
273             raise ValueError(
274                 f"indexing.chunk_overlap ({self.chunk_overlap}) must be < chunk_duration
275                 "
276                 f"({self.chunk_duration}); otherwise chunking never advances."
277             )
278         return self
279
280     class OutputConfig(BaseModel):
281         default_highlights_name: str = "highlights.mp4"
282         db_name: str = "sports_indexer.db"
283         # Directory where the DB, highlight reels, and processed clips are written.
284         # The CLI's --output-dir overrides this at startup (see backend/main.py:main).
285         output_dir: str = "output"
286
287     class WatermarkConfig(BaseModel):
288         """Branding overlay burned into each highlight clip during the per-clip re-encode
289         (backend/pipeline/stitcher/compiler.py). **On by default** ? set `enabled: false`
290         (under `stitching.watermark` in config.json) to turn it off. The text is fully
291         configurable. The concat stage stays stream-copy (-c copy)."""
292
293         enabled: bool = True
294         text: str = "Generated by Khelsutra.guru"
295         logo_path: str = "" # optional transparent PNG overlay; "" = no logo
296         font_path: str = "" # optional .ttf; "" = use the fontconfig `font` family below
297         font: str = "Sans" # fontconfig family used when font_path is empty
298         font_size_ratio: float = Field(
299             default=0.035, gt=0.0, le=0.5
300         ) # text height as a fraction of frame height
301         opacity: float = Field(default=0.85, ge=0.0, le=1.0)
302         box: bool = True # faint backing box behind the text for legibility
303         margin: int = Field(default=24, ge=0) # px inset from the chosen corner
304         position: Literal["bottom-right", "bottom-left", "top-right", "top-left"] = (
305             "bottom-right"
306         )
307
308     class StitchingConfig(BaseModel):
309         begin_padding: float = 0.5
310         end_padding: float = 0.5
311         use_cache: bool = False
312         min_shot_count: int = 1
313         # Long-rally reel filter (seconds; 0 = OFF). Drop detected rally pairs whose
314         duration
315         # (end ? start) is below this from the highlight reel, BEFORE compiling. A cleaner
316         # reel-selection knob than `min_shot_count`: shot_count is unlabeled / "Unknown" on
317         the
318         # fully-local no-AI path, whereas duration is derived from the rally's own
319         start/end. The
320         # local detector over-fires and its false positives are mostly SHORT (motion blips,
321         # background-court fragments), so this keeps the reel to the eye-catching long
322         rallies.
323         # OFF by default ? the detector output is unchanged, only which rallies reach the
324         reel.
325         min_rally_duration: float = Field(default=0.0, ge=0.0)
326         watermark: WatermarkConfig = Field(default_factory=WatermarkConfig)
327
328     class LoggingConfig(BaseModel):
329         """Logging knobs for the run entrypoints (see backend/utils/logging_setup.py)."""
330
331         # Third-party HTTP/RPC client loggers (httpx, httpcore, urllib3, google-genai,

```

```

330     # google.auth, grpc) emit one INFO line per request ? dozens per Gemini run, which
331     # buries our own [rally]/[compile] lines in run.log during manual QA. Default raises
332     # them to WARNING; set true to restore full chatter for request-flow debugging.
333     verbose_http: bool = False
334
335
336     class SecurityConfig(BaseModel):
337         # When set, /api/process video_path must resolve under this directory (hosted/tenant
338         # mode). Default None preserves the single-user local 'any local file' workflow.
339         allowed_video_root: Optional[str] = None
340
341         # --- Chunked upload (B2, PR-2) ---
342         # Landing zone for browser uploads (assembled from parts). Per-user subdirectories
343         # arrive with PR-3 identity scoping. ? In hosted mode, set allowed_video_root to
344         # contain upload_dir or /api/process will reject the uploaded paths.
345         upload_dir: str = "output/uploads"
346         # Whole-file cap (default 4 GB) and per-part cap. Parts stay under ~100 MB because
347         # free-tier edge proxies cap request bodies there (HOSTING_PLAN $2).
348         max_upload_mb: int = 4096
349         max_part_mb: int = 95
350         allowed_upload_extensions: List[str] = ["mp4", "mov"]
351
352         # --- M9 edge auth (Cloudflare Access, D9) ---
353         # DEFAULT-OFF: edge_auth_enabled=False ? no JWT is required/verified (the local
single-user
354         # path is unchanged). When True, /api/process requires + verifies the `Cf-Access-
Jwt-Assertion`
355         # header against the team JWKS (so a direct hit to the Cloud Run URL can't spoof the
identity
356         # header) and tags the job with the verified user_id (owner_id). The team domain
(e.g.
357         # ``https://<team>.cloudflareaccess.com``) supplies the JWKS + issuer; the AUD is
the CF Access
358         # application's audience tag. NO secrets here ? these are public identifiers.
359         edge_auth_enabled: bool = False
360         cf_team_domain: str = (
361             "" # e.g. https://<team>.cloudflareaccess.com (issuer + /certs JWKS)
362         )
363         cf_access_aud: str = "" # the CF Access application AUD tag the token must carry
364         # (CORS allow-origins is set via the RALLY_CORS_ALLOW_ORIGINS env var ? read at
import in
365         # backend/main.py so importing the app does no filesystem I/O; default ["*"].
Credentials stay
366         # OFF ? Cf-Access is a header, not a cookie.)
367
368
369     class StorageConfig(BaseModel):
370         """Pluggable storage backend (docs/DECOUPLED_COMPUTE/). Default ``local`` = today's
371         behaviour (filesystem paths, byte-for-byte). ``s3`` covers AWS + Cloudflare R2 + D0
Spaces
372         + MinIO (via ``endpoint_url``); ``gcs`` = Google Cloud Storage; ``gdrive`` = a
locally MOUNTED
373         Google Drive (``gdrive.mount_root``). Per-backend settings are loose dicts passed to
the constructor ?
374         **secrets are never stored here**, only endpoints / regions / file paths / source-
selectors
375         (boto3/google auth chains supply credentials). See ``backend/storage/``. """
376
377         backend: Literal["local", "s3", "gcs", "gdrive", "gdrive-api"] = "local"
378         # Output root for produced artifacts. "" => fall back to output.output_dir (non-
breaking).
379         output_root: str = ""
380         # --- Input side (COMPUTE_DECOUPLED_SERVING M2; parallel to the output
`backend`/`output_root`).
381         # Default ``local`` / "" keeps today's behaviour: a bare path / ``local://`` URI is

```

```

read IN PLACE
382      # (identity, no copy). Set these to read inputs from a bucket / mounted Drive ? a
bare/relative
383      # input name is resolved against ``input_root`` (e.g.
``input_root``="gs://bucket/videos"`` + "x.mp4"
384      # => ``gs://bucket/videos/x.mp4``) or, if ``input_root`` is empty, prefixed with
``input_backend``
385      # (``"gcs"`` + "bucket/x.mp4" => ``gcs://bucket/x.mp4``). An explicit scheme on the
input
386      # (``gs://?``/``s3://?``) or an absolute local path ALWAYS wins. Non-local inputs
are fetched to a
387      # scratch dir before processing (see ``backend.storage.acquire_local_input``).
388      input_backend: Literal["local", "s3", "gcs", "gdrive", "gdrive-api"] = "local"
389      input_root: str = ""
390      # --- Per-video workspace (docs/PER_VIDEO_WORKSPACE.md) ---
391      # One folder per video keyed by the MD5 video_id:
<root>/[<workspace_prefix>]/<video_id>/...
392      # so a local deploy and a cloud (bucket) deploy share ONE relative layout (only the
root
393      # differs). DEFAULT-OFF: convergence is phased + flipped via a Launch Report (D5).
When OFF
394      # the legacy flat layout is used unchanged.
395      single_folder_per_video: bool = False
396      # Durable workspace root URI. "" => precedence: output_root, then output.output_dir.
397      workspace_root: str = ""
398      # Cloud namespace prefix under the root so per-video folders sit tidily beside
399      # videos/ labels/ weights/ (e.g. gs://bucket/workspaces/<video_id>/). "" = root
directly.
400      workspace_prefix: str = "workspaces"
401      local: dict[str, Any] = Field(default_factory=dict)
402      s3: dict[str, Any] = Field(default_factory=dict)
403      gcs: dict[str, Any] = Field(default_factory=dict)
404      gdrive: dict[str, Any] = Field(default_factory=dict)
405      # gdrive-api settings (e.g. {"root_folder_id": "<id>"}); read via the
hyphen?underscore lookup in
406      # backend.storage.get_storage. NO secrets ? creds come from ADC over the Drive
scope.
407      gdrive_api: dict[str, Any] = Field(default_factory=dict)
408
409
410      class DeploymentConfig(BaseModel):
411          """Deployment profile ? one switch that selects a coherent bundle of compute +
storage
412          knobs (COMPUTE_DECOUPLED_SERVING M1, ``docs/COMPUTE_DECOUPLED_SERVING/README.md``
§4.3).
413
414          **Pure / additive:** the empty default profile (``""``) applies NO preset, so the
pipeline
415          behaves exactly as before this section existed. A profile is opt-in (``config.json``
416          ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); auto-detect only ever
417          *suggests* one (D15 ? cloud spend is never entered silently). Nothing in the core
reads
418          these fields yet (M1): selecting a profile works purely by pre-setting the EXISTING
knobs
419          (``indexing.detector_impl`` / ``skip_ai_handoff``, ``storage.backend``) in the
loader.
420
421          Valid values mirror ``backend/config/presets.py`` (parity-pinned by a unit test)."""
422
423          # "" = no profile (today's manual knobs). Others apply presets.PROFILE_PRESETS.
424          profile: Literal["", "truly-local", "cloud-serving", "cpu-mock"] = ""
425          # Where the core runs. Set by the profile preset; overridable in config.json. Inert
in M1
426          # (recorded for observability; consumed by the ComputeBackend seam from M2+).
427          compute_target: Literal["", "local_gpu", "cloud_run", "cpu_mock"] = ""

```

```

428
429
430     class BillingConfig(BaseModel):
431         """Per-job cost ledger (COMPUTE_DECOUPLED_SERVING M8, D8).
432
433         **DEFAULT-OFF:** ``enabled=False`` ? nothing records cost (the v6 cost columns stay
434         NULL =
435         today's behaviour). When enabled, a job's cost is computed + stored in **USD** via
436         the versioned
437         ``pricebook`` DB table and **displayed in INR** at ``fx_inr_per_usd``. The default
438         pricebook
439         version is the seeded ``placeholder-v0`` (every rate ``0.0`` ? cost ``0``) until the
440         owner
441         inserts a real, calibrated version and points ``pricebook_version`` at it.
442         ? Even with ``enabled=True`` + a real pricebook, recorded cost stays **0** until a
443         follow-up
444         wires the worker to emit per-job ``usage`` (gpu/wall/egress/gemini) ? M8 lands the
445         ledger +
446         pricebook + the recording seam; usage emission (the metering hooks) is a separate
447         step."""
448         enabled: bool = False
449         pricebook_version: str = (
450             "placeholder-v0" # the seeded $0 version; owner flips to a real one
451         )
452         fx_inr_per_usd: float = 83.0 # USD?INR display rate (Bangalore market)
453         margin: float = 0.0 # resale markup (?0); 0 = bill at cost
454
455     class FlywheelConfig(BaseModel):
456         """Data-flywheel harness (COMPUTE_DECOUPLED_SERVING M11, D12) ? the
457         consent?capture?eval?goldenize
458         loop: on a CONSENTED adult job, durably **capture** the source + rally output into
459         the per-video
460         workspace ? **shadow-eval** owned-vs-Gemini ? **propose promotion** to the golden
461         TRAINING set, so
462         the owned model climbs toward the D11 ship floor.
463
464         **DEFAULT-OFF + INERT:** ``enabled=False`` ? nothing is captured/evaluated (byte-
465         identical to
466         today). ? **Consent is FAKED-DENY until counsel + the consent schema land:** even
467         ``enabled=True``
468         captures **nothing**, because ``flywheel.consent_attested`` currently returns
469         ``False`` for every
470         job (the privacy guardrail ? D10: adults-only, explicit opt-in, derived-data + auto-
471         delete-raw).
472         This is the PLUMBING only; the real consent gate (counsel-reviewed ToS/DPA ? consent
473         columns ?
474         wiring) + the owned-model + golden-label data needed to *activate* shadow-
475         eval/promotion are
476         owner/counsel/GPU scope. Pairs with M9-consent."""
477         enabled: bool = False
478         workspace_prefix: str = (
479             "flywheel" # per-video workspace prefix for captured artifacts
480         )
481         shadow_eval: bool = True # run the owned-vs-Gemini shadow comparison on capture
482         min_promote_n: int = 5 # promote_if_served_safe `min_n` (the served-safe gate)
483
484     class AppConfig(BaseModel):
485         """Root configuration object.
486
487         All runtime configuration should be accessed via an instance of this class

```



```

477         (or the loader) rather than raw dict lookups.
478         """
479
480         sport: Sport = "badminton"
481         ai_provider: str = "gemini"
482         # Auto-tracking alias: 2.5-flash is legacy now that 3.5-flash is the stable
flagship;
483         # the alias never deprecates underneath us. Eval reports record the run date
alongside.
484         ai_model: str = "gemini-flash-latest"
485
486         validation: ValidationConfig = Field(default_factory=ValidationConfig)
487         indexing: IndexingConfig = Field(default_factory=IndexingConfig)
488         output: OutputConfig = Field(default_factory=OutputConfig)
489         stitching: StitchingConfig = Field(default_factory=StitchingConfig)
490         security: SecurityConfig = Field(default_factory=SecurityConfig)
491         logging: LoggingConfig = Field(default_factory=LoggingConfig)
492         storage: StorageConfig = Field(default_factory=StorageConfig)
493         deployment: DeploymentConfig = Field(default_factory=DeploymentConfig)
494         billing: BillingConfig = Field(default_factory=BillingConfig)
495         flywheel: FlywheelConfig = Field(default_factory=FlywheelConfig)
496
497         # Reject unknown keys ? typed config migration is complete (D3).
498         model_config = {
499             "extra": "forbid",
500             "validate_assignment": True,
501         }
502
503
504         # Convenience type alias for callers that want to annotate
505         Config = AppConfig
506

```

8. assistant (2026-07-03T19:48:37.076Z)

Let me find the presets file and check the WASB tracker implementation to understand `tracker.update()` cost. Let me also look at how `det\_by\_fid` is built to reason about the tracker loop.

9. user (2026-07-03T19:48:37.660Z)

No files found

10. user (2026-07-03T19:48:41.321Z)

```

/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/backend/config/loader.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/backend/config/models.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/backend/config/presets.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/backend/config/startup.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/backend/config/__init__.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/backend/config/__pycache__/loader.cpython-313.pyc
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/backend/config/__pycache__/models.cpython-313.pyc
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/backend/config/__pycache__/presets.cpython-313.pyc

```

```
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/backend/config/__pycache__/startup.cpython-313.pyc
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/backend/config/__pycache__/__init__.cpython-313.pyc
```

11. assistant (2026-07-03T19:48:44.754Z)

The Glob path issue is because the working dir differs from the worktree. Let me read presets.py and search for the WASB tracker source. The tracker itself lives at /opt/models/WASB-SBDT which I can't read, but let me check if there's a vendored copy or reference anywhere, and read presets.

12. user (2026-07-03T19:48:46.292Z)

```
1      """Deployment-profile presets + environment auto-detection (COMPUTE_DECOUPLED_SERVING
M1).
2
3      A *deployment profile* is a single switch that applies a coherent BUNDLE of existing
4      config knobs, so the truly-local / cloud-serving / cpu-mock modes don't drift apart one
5      knob at a time (the design in ``docs/COMPUTE_DECOUPLED_SERVING/README.md`` §4.3).
6
7      Precedence (lowest ? highest), realised by the loader::
8
9          built-in model defaults < profile preset < config.json < env < worker --set
10
11      This module is **pure data + pure functions** ? no IO, and (critically) **no torch
import at
12      module load**. The loader applies a preset ONLY when a profile is EXPLICITLY selected
13      (``config.json`` ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); with no
14      profile selected the loader is byte-identical to the pre-M1 behaviour (``profile=""`` ==
15      today). Auto-detect only ever **suggests** a profile ? it never auto-applies one (D15:
cloud
16      spend is never entered silently).
17      """
18
19      from __future__ import annotations
20
21      import copy
22      import os
23      from typing import Any, Mapping, Optional
24
25      # Canonical enum values, defined ONCE here. ``models.DeploymentConfig``'s ``Literal``
mirrors
26      # these and ``tests/test_config_deployment.py`` pins parity so the two can never drift.
27      PROFILES: tuple[str, ...] = ("truly-local", "cloud-serving", "cpu-mock")
28      COMPUTE_TARGETS: tuple[str, ...] = ("local_gpu", "cloud_run", "cpu_mock")
29
30      # Each profile maps 1:1 to its canonical compute_target. Used to apply the preset and to
warn
31      # when config.json explicitly overrides compute_target into a state inconsistent with
the profile.
32      COMPUTE_TARGET_FOR_PROFILE: dict[str, str] = {
33          "truly-local": "local_gpu",
34          "cloud-serving": "cloud_run",
35          "cpu-mock": "cpu_mock",
36      }
37
38      # A preset is a partial, nested override dict over the model defaults. It touches ONLY
knobs
39      # that EXIST today ? M1 is pure/additive: the new
``deployment.{profile,compute_target}`` plus
40      # ``indexing.{detector_impl,skip_ai_handoff}`` and ``storage.backend``. Later milestones
extend
```

```

41 # these bundles ? ``input_backend`` (M2) and the configurable output/workspace base (M3,
42 # Launch-Report-gated) ? deliberately NOT set here so M1 never pre-empts a gated flip.
43 PROFILE_PRESETS: dict[str, dict[str, Any]] = {
44     "truly-local": {
45         "deployment": {"profile": "truly-local", "compute_target": "local_gpu"},
46         # detector_impl stays "auto" (the runner picks WSL vs native by env on the GPU
47         box).
48         # Zero-cloud / privacy-first ? run fully local with no Gemini handoff.
49         "indexing": {"skip_ai_handoff": True},
50         "storage": {"backend": "local"},
51     },
52     "cloud-serving": {
53         "deployment": {"profile": "cloud-serving", "compute_target": "cloud_run"},
54         # Linux L4 (cuda); Gemini handoff ON until the owned model clears the #172 floor
55         (D11).
56         # wasb GPU-feed tuning: the parity defaults (batch 8, no prefetch) left the paid
57         L4
58         # <1% utilized on the first real serving run (2026-07-03) ? the pipeline was
59         # CPU-feed-bound. Serving opts into the measured fast values; local/WSL
60         unchanged.
61         "indexing": {
62             "detector_impl": "native",
63             "skip_ai_handoff": False,
64             "wasb": {"batch_size": 64, "prefetch_batches": 4},
65         },
66         "storage": {"backend": "gcs"},
67     },
68     "cpu-mock": {
69         "deployment": {"profile": "cpu-mock", "compute_target": "cpu_mock"},
70         # Deterministic CPU mock ? no GPU/WSL, no provider/API. The CI / regression
71         profile.
72         "indexing": {"detector_impl": "stub", "skip_ai_handoff": True},
73         "storage": {"backend": "local"},
74     },
75 }
76
77 def is_known_profile(profile: str) -> bool:
78     """True for a non-empty, recognised profile name."""
79     return bool(profile) and profile in PROFILE_PRESETS
80
81 def preset_for(profile: str) -> dict[str, Any]:
82     """Return a deep COPY of the preset bundle for ``profile`` (so callers can't mutate
83     the
84     module-level table), or ``{}`` for ``""`` / an unknown profile."""
85     return copy.deepcopy(PROFILE_PRESETS.get(profile, {}))
86
87 def _truthy(val: Optional[str]) -> bool:
88     """Common env-var truthiness (``1/true/yes/on``, case/space-insensitive)."""
89     if val is None:
90         return False
91     return val.strip().lower() in {"1", "true", "yes", "on"}
92
93 def _cuda_available() -> bool:
94     """Best-effort CUDA probe. torch is heavy + optional, so it is imported lazily here
95     and
96     ONLY when a caller opts into ``probe_cuda``; any import/runtime failure means 'no
97     GPU'."""
98     try:
99         # Intentionally lazy + local: keeps the config-load path torch-free (torch is
100         heavy
101         # and optional). Only reached when a caller opts into probe_cuda.

```

```

97         import torch
98
99         return bool(torch.cuda.is_available())
100     except Exception:
101         return False
102
103
104     def probe_cuda_available() -> bool:
105         """Public best-effort CUDA probe (lazy torch import; ``False`` on any failure).
106
107         Exposed for the per-profile startup checks (``backend/config/startup.py``) so a
108         caller that *wants* the GPU heads-up can pass the result in without the startup module
109         importing torch itself. Best-effort by design: if torch lives only in a detector subprocess
110         env (native WASB), this returns ``False`` in the main process ? which is why the startup
111         checks treat a GPU mismatch as a WARNING (the runner healthcheck is the authority), never a
112         hard error that could false-fail."""
113         return _cuda_available()
114
115
116     def suggest_profile(
117         env: Optional[Mapping[str, str]] = None, *, probe_cuda: bool = False
118     ) -> str:
119         """SUGGEST (never auto-apply ? D15) a deployment profile from the environment.
120
121         Signals, in priority order:
122
123         - ``K_SERVICE`` set (Cloud Run)           ? ``"cloud-serving"``
124         - ``CI`` truthy                             ? ``"cpu-mock"``
125         - ``probe_cuda`` and a CUDA GPU present ? ``"truly-local"``
126         - otherwise                               ? ``""`` (no strong suggestion)
127
128         ``probe_cuda`` defaults to ``False`` so the common config-load path stays cheap and
129         never imports torch; a CLI ``--detect`` or cloud-serving startup that wants the GPU probe
130         passes ``probe_cuda=True``. The result is advisory only: the caller decides whether to act
131         on it.
132
133         """
134         e = os.environ if env is None else env
135         if e.get("K_SERVICE"):
136             return "cloud-serving"
137         if _truthy(e.get("CI")):
138             return "cpu-mock"
139         if probe_cuda and _cuda_available():
140             return "truly-local"
141         return ""

```

13. user (2026-07-03T19:48:48.431Z)

(Bash completed with no output)

14. assistant (2026-07-03T19:48:50.941Z)

Now let me look at the WASB tracker source to understand the cost of `tracker.update()`. Let me check the Dockerfile and search for any vendored WASB tracker code or references to the online tracker algorithm.

15. user (2026-07-03T19:48:53.434Z)

```

1      # Cloud Run GPU **Job** image for the rally worker (COMPUTE_DECOUPLED_SERVING M6/M7,
ADR-012/D16).
2      #
3      # Runs `python -m backend.worker infer ?` INLINE on an L4 (detect + window + stitch co-
located so
4      # CPU/wall + egress are metered into one job ? D4). Built + pushed + verified by the
owner on the
5      # M7 real run (the $100/?10k budget, D17); CI does NOT build this (a CUDA image is too
heavy) ? the
6      # dispatch shim is proven with a fake client (tests/test_compute_backend.py).
7      #
8      # Two python stacks, deliberately separate (a framework conflict otherwise ? same split
as the
9      # native GPU box, deploy/digitalocean/gpu_setup.sh):
10     # * the MAIN app env (this repo + ffmpeg) drives the pipeline + stitching;
11     # * a conda `wasb` env (py3.8 + torch 1.11.0+cu113) runs the WASB shuttle detector as
a subprocess
12     # (the native runner shells out to $WASB_PYTHON), so the app's torch never co-
installs with it.
13     #
14     # ? WASB-C3: the pretrained badminton weights are academic-data-trained ? provenance NOT
cleared for
15     # commercial sharing. They are NOT baked into this image; stage them at runtime (a
bucket fetch or a
16     # mount) and point $WASB_WEIGHTS at them. Resolve WASB-C3 before any public, non-owner
serving.
17
18     FROM nvidia/cuda:12.4.1-runtime-ubuntu22.04
19
20     ENV DEBIAN_FRONTEND=noninteractive \
21         PYTHONUNBUFFERED=1 \
22         PIP_NO_CACHE_DIR=1 \
23         CONDA_DIR=/opt/conda \
24         WASB_REPO_DIR=/opt/models/WASB-SBDT
25
26     # ffmpeg (stitch/decode) + DejaVu fonts (the watermark filter) + git/curl for the WASB
env build.
27     RUN apt-get update && apt-get install -y --no-install-recommends \
28         ffmpeg fonts-dejavu-core git curl ca-certificates build-essential \
29         python3 python3-pip python3-venv \
30         && rm -rf /var/lib/apt/lists/*
31
32     # --- WASB detector env (micromamba + conda-forge ? BSD-licensed, NO Anaconda ToS; P3a,
ADR-005) ---
33     # Enforces ADR-005 (every dependency ? code, data, weights, AND the build toolchain ?
stays
34     # MIT/Apache/BSD): micromamba (BSD-3) + conda-forge (community) replaces Miniconda's
Anaconda-ToS-gated
35     # `defaults` channels (the ?200-employee Terms-of-Service trap flagged in
PACKAGING_PLAN.md §3).
36     # MAMBA_ROOT_PREFIX=$CONDA_DIR (= /opt/conda) keeps the env at $CONDA_DIR/envs/wasb so
$WASB_PYTHON below
37     # is UNCHANGED. Same verified torch-1.11.0+cu113 WASB stack; the pip wheels are
unaffected by the switch.
38     ENV MAMBA_ROOT_PREFIX=/opt/conda
39     RUN curl -sSL https://micro.mamba.pm/api/micromamba/linux-64/latest \
40         | tar -xj -C /usr/local bin/micromamba \
41         && micromamba create -y -q -n wasb -c conda-forge python=3.8 \
42         && git clone --depth 1 https://github.com/nttcom/WASB-SBDT.git "$WASB_REPO_DIR" \
43         && "$CONDA_DIR/envs/wasb/bin/python" -m pip install -q \
44             torch==1.11.0+cu113 torchvision==0.12.0+cu113 --extra-index-url
https://download.pytorch.org/whl/cu113 \
45         && ( [ -f "$WASB_REPO_DIR/requirements.txt" ] && "$CONDA_DIR/envs/wasb/bin/python"
-m pip install -q -r "$WASB_REPO_DIR/requirements.txt" || true ) \
46         && "$CONDA_DIR/envs/wasb/bin/python" -m pip install -q \

```

```

47         hydra-core omegaconf pandas opencv-python-headless pillow scipy tqdm pyyaml
matplotlib "numpy<1.24"
48
49     # --- Main application env ---
50     WORKDIR /app
51     COPY pyproject.toml README.md ./
52     COPY requirements-trackers.txt ./
53     COPY backend ./backend
54     COPY training ./training
55     # Install with the [storage-gcs] extra: the cloud worker PULLS the gs:// input and
PUSHES the
56     # gs:// outputs/done-marker through backend.storage (google-cloud-storage is lazy-
imported there).
57     # Without it the job dies at the first gs:// fetch with `ModuleNotFoundError:
google.cloud` ? the
58     # real-Cloud-Run failure the mock-client tests can't see (the bucket I/O is faked in
tests).
59     # D13/P12: `trackers` (ByteTrack's new home) is a git-only dep (broken PyPI wheel ?
60     # roboflow/trackers#474) so it's NOT pulled by `pip install .`; install it out-of-band
before
61     # the main env, or the person-cue path crashes with ModuleNotFoundError. The immutable
git SHA is
62     # single-sourced in requirements-trackers.txt.
63     RUN python3 -m pip install --upgrade pip \
64         && python3 -m pip install -r requirements-trackers.txt \
65         && python3 -m pip install ".[storage-gcs]"
66
67     # The native runner finds the WASB env + repo by env (no `conda activate` needed);
weights come in at
68     # runtime (WASB-C3). detector_impl=native = the Linux-native WASB runner; cuda stays the
default.
69     # PYTHONUNBUFFERED: without it the worker's stdout/stderr sit in the block buffer until
process
70     # exit, so Cloud Logging shows a single end-of-run burst instead of live progress
(observed on the
71     # first real M7 runs, 2026-07-03) ? a running job looks stalled and mid-run failures log
nothing.
72     ENV PYTHONUNBUFFERED=1 \
73         WASB_PYTHON=/opt/conda/envs/wasb/bin/python \
74         WASB_REPO_DIR=/opt/models/WASB-SBDT \
75         WASB_WEIGHTS=/opt/models/WASB-SBDT/pretrained_weights/wasb_badminton_best.pth.tar \
76         RALLY_DETECTOR_IMPL=native
77
78     # The dispatcher (CloudRunCompute) supplies the per-execution args:
79     # infer --video-uri gs:/// --out-uri gs:/// --done-uri gs:/// --segmenter wasb_hybrid
--set ?
80     ENTRYPOINT ["python3", "-m", "backend.worker"]
81

```

16. user (2026-07-03T19:48:59.796Z)

```

---refresh---
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.github/workflows/nightly.yml:230: # failure, open/refresh a single tracking issue so
a human is pinged. Uses the
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.github/workflows/nightly.yml:235: name: Nightly · Surface failure (open/refresh
tracking issue)
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-
highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.0.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-
highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-

```

```

worktree/.mypy_cache/3.13/cache.10.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.12.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.13.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.14.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.15.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.2.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.3.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.4.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.5.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.6.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.7.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.8.db matches
Binary file /c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.mypy_cache/3.13/cache.9.db matches
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.pytest_cache/v/cache/nodeids:744:
"tests/test_golden_real_fixtures.py::test_refresh_snapshot_cli_requires_reason",
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/.pytest_cache/v/cache/nodeids:745:
"tests/test_golden_real_fixtures.py::test_refresh_snapshot_no_target_raises",
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/backend/eval/calibrate_local.py:119: of times, so the floor refresh uses this cached table instead.
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/backend/eval/golden_fixtures.py:421:# Freeze / refresh (expected.json + provenance.json are correct-by-construction)

```

17. assistant (2026-07-03T19:49:04.978Z)

The WASB tracker source isn't vendored in the repo (it lives in ``/opt/models/WASB-SBDT`` which I can't read). I have enough to reason from the actual code in ``wasb_infer.py``. Let me verify one crucial detail about the tracker loop's cost driver ? how large ``det_by_fid[fid]`` can get. The key concern is ``_dets_to_tracker`` which allocates numpy arrays per candidate per frame. Let me also check what the WASB online tracker's ``update()`` does algorithmically. Let me search the WASB repo reference more broadly and confirm the candidate-count characteristics.

18. user (2026-07-03T19:49:09.226Z)

```

=== wasb_runner exists? ===
total 196
drwxr-xr-x 1 avidu 197609      0 Jul  4 01:13 .
drwxr-xr-x 1 avidu 197609      0 Jul  4 00:49 ..
-rw-r--r-- 1 avidu 197609 1065 Jul  4 00:49 __init__.py
drwxr-xr-x 1 avidu 197609      0 Jul  4 01:16 __pycache__
-rw-r--r-- 1 avidu 197609 1666 Jul  4 00:49 audio_features.py
-rw-r--r-- 1 avidu 197609 11435 Jul  4 00:49 base.py
-rw-r--r-- 1 avidu 197609 5604 Jul  4 00:49 court_gate.py
-rw-r--r-- 1 avidu 197609 2493 Jul  4 00:49 device.py
-rw-r--r-- 1 avidu 197609 10063 Jul  4 00:49 fusion_features.py
-rw-r--r-- 1 avidu 197609 16761 Jul  4 01:11 native_wasb_runner.py
-rw-r--r-- 1 avidu 197609 19923 Jul  4 00:49 person_detector.py
-rw-r--r-- 1 avidu 197609 5447 Jul  4 00:49 rally_gate.py
-rw-r--r-- 1 avidu 197609 8269 Jul  4 00:49 stub_runner.py
-rw-r--r-- 1 avidu 197609 28469 Jul  4 00:49 tracknet_runner.py
-rw-r--r-- 1 avidu 197609 39094 Jul  4 01:13 wasb_infer.py
-rw-r--r-- 1 avidu 197609 10610 Jul  4 00:49 wasb_runner.py

```

19. assistant (2026-07-03T19:49:13.907Z)

There's no vendored WASB tracker (it's in the conda env). I have enough from the real `wasb\_infer.py` code to reason precisely. Let me confirm the total frame-count scale for a 7.5-min 60fps clip and characterize the loop cost. Let me check the parity test that guards output determinism, since parity_risk is load-bearing.

20. user (2026-07-03T19:49:16.334Z)

```

/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_ablation.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_annotation_proxy.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_api_cloud_dispatch.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_api_endpoints.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_app_home.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_async_jobs.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_compute_backend.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_config_deployment.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_database.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_eval_stats.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_experiment.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_ffmpeg_resolver.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_fusion_golden_telemetry.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_golden_fixtures.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-

```



```
worktree/tests/test_golden_real_fixtures.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/tests/test_job_worker_loop.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/tests/test_motion_segmenter.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/tests/test_native_wasb_runner.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/tests/test_nightly_reelcount.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/tests/test_nightly_regression.py
```

21. assistant (2026-07-03T19:49:19.356Z)

Let me look at the streaming-vs-disk parity test specifically, to be precise about what the byte-parity guarantee actually asserts (this determines parity_risk for any tracker-loop change).

22. user (2026-07-03T19:49:23.330Z)

```
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/tests/test_native_wasb_runner.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-worktree/tests/test_wasb_cache.py
/c/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/fix-
worktree/tests/___pycache___/test_wasb_cache.cpython-313-pytest-9.0.3.pyc
```

23. user (2026-07-03T19:49:27.980Z)

```
37: def test_det_plain_and_back_roundtrip():
42:     back = wi._dets_to_tracker([plain])
52: def test_replay_accumulates_detections_across_windows(tmp_path):
55:     # frame 1 appears in BOTH windows -> its candidates must accumulate in window order;
72: def test_resume_index_equals_window_count(tmp_path):
80: def test_empty_cache(tmp_path):
90: def test_truncated_trailing_line_recovered_and_cleaned(tmp_path):
102:     # file is rewritten to the clean prefix -> appending window 2 stays contiguous
110: def test_noncontiguous_window_index_stops_replay(tmp_path):
124: def test_manifest_roundtrip_and_compat(tmp_path):
158: def test_manifest_device_keying_back_compat():
189: def test_manifest_compatible_streaming_video_key():
233: def test_read_manifest_missing_returns_none(tmp_path):
237: def test_reset_cache_removes_files(tmp_path):
250: def test_default_cache_dir_is_next_to_output():
257: def test_atomic_write_trajectory(tmp_path):
271: def test_synthetic_frame_path_roundtrip():
281: def test_build_windows_matches_disk_logic():
291: def test_build_windows_too_few_frames_is_empty():
299:     """A fake forward-only decoder: returns ('frame', i) and records call order."""
309: def test_store_reads_each_frame_once_in_order_for_sliding_windows():
310:     # The detector's real access pattern: ImageDataset reads window i's frames in order
316:     order = [0, 1, 2, 1, 2, 3, 2, 3, 4]
317:     got = [store.get(i) for i in order]
318:     assert got == [("frame", i) for i in order]
319:     assert calls == [0, 1, 2, 3, 4] # decoded once each, in order
322: def test_store_evicts_outside_window_bounded_memory():
334: def test_store_keeps_window_minus_one_lookback():
```

```

345: def test_store_resume_starts_at_start_index_without_decoding_prefix():
356: def test_store_out_of_range_raises():
374: def test_make_streamed_read_image_returns_pil_rgb_and_falls_through():
400: def test_make_streamed_read_image_resume_primes_start_index():
414: def test_bgr_to_pil_rgb_matches_disk_roundtrip(tmp_path):

```

24. assistant (2026-07-03T19:49:30.687Z)

Let me look at the key test that establishes the ordering contract for `det\_by\_fid`, which is exactly what any vectorization must preserve.

25. user (2026-07-03T19:49:33.741Z)

```

1      """Resumable-cache layer of wasb_infer (pure-Python parts; no WSL/GPU).
2
3      Covers the reboot-resume guarantees: per-window incremental cache, cross-window
4      detection accumulation on replay, crash-truncation recovery, manifest
5      compatibility/invalidation, and atomic trajectory writes.
6
7      Also covers the streaming fast-path's pure helpers (synthetic frame addressing,
8      window builder, the sliding-buffer frame store, and the read_image shim) ? these
9      let the detector read frames straight from the decoded video instead of from ~46k
10     PNGs on disk, while preserving output exactly (byte-identical, verified).
11     """
12
13     import json
14     import os
15     import os.path as osp
16
17     import numpy as np
18     import pytest
19
20     from backend.pipeline.detectors import wasb_infer as wi
21
22
23     def _win(idx, frames):
24         """A window record: frames = [(fid, [[x,y,score,scale], ...]), ...]."""
25         return {"w": idx, "f": [[fid, dets] for fid, dets in frames]}
26
27
28     def _write_windows(cache_dir, windows):
29         det_jsonl, _ = wi._cache_paths(cache_dir)
30         with open(det_jsonl, "a") as f:
31             wi._append_windows(f, windows)
32
33
34     # ----- #
35     # detection (de)serialization
36     # ----- #
37     def test_det_plain_and_back_roundtrip():
38         det = {"xy": np.array([439.0, 64.5]), "score": np.float32(0.87), "scale": 0}
39         plain = wi._det_plain(det)
40         assert plain == [439.0, 64.5, float(np.float32(0.87)), 0]
41
42         back = wi._dets_to_tracker([plain])
43         assert len(back) == 1
44         assert isinstance(back[0]["xy"], np.ndarray) # tracker does vector math on xy
45         assert back[0]["xy"].tolist() == [439.0, 64.5]
46         assert back[0]["scale"] == 0
47
48
49     # ----- #
50     # cache replay + resume index
51     # ----- #
52     def test_replay_accumulates_detections_across_windows(tmp_path):

```

```

53     cache = str(tmp_path / "c")
54     os.makedirs(cache)
55     # frame 1 appears in BOTH windows -> its candidates must accumulate in window order;
56     # frame 2 is empty in window 0 but detected in window 1.
57     _write_windows(
58         cache,
59         [
60             _win(0, [(0, [[10, 20, 0.9, 0]]), (1, [[11, 21, 0.8, 0]]), (2, [])]),
61             _win(1, [(1, [[12, 22, 0.7, 0]]), (2, [[13, 23, 0.6, 0]]), (3, [])]),
62         ],
63     )
64     windows_done, det = wi.load_and_clean_cache(cache)
65     assert windows_done == 2
66     assert det[0] == [[10, 20, 0.9, 0]]
67     assert det[1] == [[11, 21, 0.8, 0], [12, 22, 0.7, 0]] # window-0 then window-1
68     assert det[2] == [[13, 23, 0.6, 0]]
69     assert det[3] == [] # seen but never detected
70
71
72     def test_resume_index_equals_window_count(tmp_path):
73         cache = str(tmp_path / "c")
74         os.makedirs(cache)
75         _write_windows(cache, [_win(i, [(i, [[i, i, 0.5, 0]])]) for i in range(5)])

```

26. assistant (2026-07-03T19:49:41.093Z)

I now have a complete picture. Let me do the quantitative analysis of the tracker-postproc loop.

Key facts established from the real code:

****The tracker re-run loop**** (`wasb\_infer.py:684-690`):

```

```python
tracker = build_tracker(cfg)
tracker.refresh()
rows = []
for fid in sorted(det_by_fid.keys()):
 r = tracker.update(_dets_to_tracker(det_by_fid[fid]))
 rows.append((fid, 1 if r.get("visi") else 0, r.get("x", -1), r.get("y", -1)))
...

```

**\*\*Scale for a 7.5-min 60fps clip:\*\*** ~27,000 frames ? ~27,000 windows (`build\_windows` = `len(frames) - frames\_in + 1`), and the loop iterates `sorted(det\_by\_fid.keys())` = ~27,000 frames.

Let me verify the critical numbers on total frame count and per-frame candidate accumulation semantics, and check whether the tracker is inherently serial (it's a stateful online tracker ? `refresh()` then per-frame `update()`).

## 27. user (2026-07-03T19:49:43.416Z)

```

1 """WASB shuttle-trajectory runner (WSL2 subprocess adapter).
2
3 Windows-side counterpart to backend/detectors/wasb_infer.py. It:
4 1. stages the video into the WSL filesystem (fast local reads, not /mnt),
5 2. copies wasb_infer.py into the WASB repo's src/ dir,
6 3. runs it in the WSL `wasb` conda env -> a Frame,Visibility,X,Y trajectory CSV,
7 4. reuses the (model-agnostic) trajectory->action-window logic from tracknet_runner.
8
9 So the only WASB-specific bits live here; the windowing + CSV parsing are shared
10 with the TrackNet path. Verified inference core: wasb_infer.py (see its docstring).
11 """
12
13 import logging
14 import os
15 import shlex

```

```

16 import subprocess
17 from dataclasses import dataclass
18 from typing import Optional, Tuple
19
20 from backend.config.models import IndexingConfig
21 from backend.pipeline.detectors.base import (
22 DetectorRunner,
23 WslCommandMixin,
24 wsl_tilde_quote,
25)
26
27 # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical.
28 # parse_trajectory_csv / trajectory_to_action_windows are @staticmethods on
29 TrackNetRunner.
30 from backend.pipeline.detectors.tracknet_runner import TrackNetRunner, to_wsl_mnt_path
31
32 logger = logging.getLogger(__name__)
33
34 # Windows path to the inference wrapper we stage into WSL.
35 _INFER_WIN = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")
36
37 @dataclass
38 class WasbConfig:
39 """Resolved settings for a WASB WSL run (built from config.json ->
40 indexing.wasb)."""
41 distro: str = "Ubuntu"
42 repo_dir: str = "~/models/WASB-SBDT"
43 conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
44 conda_env: str = "wasb"
45 weights_path: str = (
46 "~/models/WASB-SBDT/pretrained_weights/wasb_badminton_best.pth.tar"
47)
48 sport: str = "badminton"
49 wsl_stage_dir: str = "~/clips"
50 timeout_sec: int = 1800 # 30 min; a hung GPU/conda call is killed (0 = no timeout)
51 keep_frames: bool = (
52 False # retain staged video + frame cache after success (debug only)
53)
54 min_free_gb: float = 0.0 # absolute floor (GB) atop the ~15x-source BLOCK guard (0
55 = off)
56 # FAST PATH (default OFF until verified): decode the staged video on the fly inside
57 # WSL and feed frames straight to the detector, skipping the slow per-frame PNG
58 # extraction that dominates a long clip. Output is bit-identical to the PNG path
59 # (PNG is lossless), but this stays opt-in until the owner's side-by-side confirms
60 it.
61 stream_video: bool = False
62
63 @classmethod
64 def from_indexing_cfg(cls, idx_cfg: "IndexingConfig") -> "WasbConfig":
65 if isinstance(idx_cfg, dict):
66 idx_cfg = IndexingConfig(**idx_cfg)
67 w = idx_cfg.wasb
68 return cls(
69 distro=w.wsl_distro,
70 repo_dir=w.repo_dir,
71 conda_sh=w.conda_sh,
72 conda_env=w.conda_env,
73 weights_path=w.weights_path,
74 sport=w.sport,
75 wsl_stage_dir=w.wsl_stage_dir,
76 timeout_sec=w.timeout_sec,
77 keep_frames=w.keep_frames,
78 min_free_gb=w.min_free_gb,

```

```

77 stream_video=w.stream_video if w.stream_video is not None else False,
78)
79
80
81 class WasbRunner(WslCommandMixin, DetectorRunner):
82 def __init__(self, cfg: WasbConfig):
83 self.cfg = cfg
84
85 def healthcheck(self) -> Tuple[bool, str]:
86 """Verify the WSL `wasb` env imports torch and sees a GPU. Returns (ok, msg)."""
87 cmd = (
88 f"conda activate {self.cfg.conda_env} && "
89 f"python -c \"import torch; print('torch', torch.__version__, "
90 f"'cuda', torch.cuda.is_available())\""
91)
92 try:
93 res = self._wsl_bash(cmd)
94 except FileNotFoundError:
95 return False, "`wsl` not found ? Windows + WSL2 required."
96 except subprocess.TimeoutExpired:
97 return False, "WSL healthcheck timed out."
98 out = (res.stdout or "").strip()
99 if res.returncode != 0:
100 return False, f"WSL `wasb` env check failed: {(res.stderr or out).strip()}"
101 return True, out
102
103 def _stage_video(self, video_win_path: str, dest_name: str) -> Optional[str]:
104 src_mnt = to_wsl_mnt_path(video_win_path)
105 reason = self.wsl_source_unreadable_reason(src_mnt)
106 if reason:
107 logger.error("Cannot stage video into WSL: %s", reason)
108 return None
109 dest = f"{self.cfg.wsl_stage_dir}/{dest_name}"
110 cmd = f"mkdir -p {self.cfg.wsl_stage_dir} && cp {shlex.quote(src_mnt)}"
111 {wsl_tilde_quote(dest)} && echo OK"
112 try:
113 res = self._wsl_bash(cmd)
114 except subprocess.TimeoutExpired:
115 logger.error("Staging video into WSL timed out.")
116 return None
117 if res.returncode != 0 or "OK" not in (res.stdout or ""):
118 logger.error(
119 "Failed to stage video into WSL: %s", (res.stderr or res.stdout).strip()
120)
121 return None
122 return dest
123
124 def _stage_infer_script(self) -> bool:
125 """Copy wasb_infer.py into the WASB repo src/ so it can import WASB modules."""
126 src_mnt = to_wsl_mnt_path(_INFER_WIN)
127 cmd = f"cp {shlex.quote(src_mnt)} {self.cfg.repo_dir}/src/wasb_infer.py && echo"
128 OK"
129 res = self._wsl_bash(cmd)
130 return res.returncode == 0 and "OK" in (res.stdout or "")
131
132 def run_predict(
133 self, video_win_path: str, output_win_dir: str, video_id: Optional[str] = None
134) -> Optional[str]:
135 """Run WASB on a video. Returns the Windows path to the trajectory CSV, or None.
136
137 All WSL/cache/output artifacts are namespaced by ``video_id`` (the file MD5) so
138 two videos that share a filename (e.g. two ``match.mp4`` from different folders)
139 can never reuse each other's staged frames, resumable cache, or CSV.
140
141 # Resolve relative dirs (the hybrids pass "output/wasb") BEFORE the /mnt

```

```

140 # translation: to_wsl_mnt_path passes relative paths through unchanged, so
141 # WSL resolved them against the inference cwd (~/.models/WASB-SBDT/src) and
142 # the CSV landed inside WSL while Windows polled the repo-relative path.
143 output_win_dir = os.path.abspath(output_win_dir)
144 os.makedirs(output_win_dir, exist_ok=True)
145 # Normalize for cross-platform basename (tests use Windows paths on Linux).
146 video_win_path = str(video_win_path).replace("\\", "/")
147 base = os.path.basename(video_win_path)
148 stem, ext = os.path.splitext(base)
149 # Unique, content-derived key: same filename + different bytes -> different key.
150 key = f"{stem}__{video_id[:12]}" if video_id else stem
151 wsl_out_dir = to_wsl_mnt_path(output_win_dir)
152 expected_csv_win = os.path.join(output_win_dir, f"{key}_wasb.csv")
153
154 if not self._stage_infer_script():
155 logger.error("Could not stage wasb_infer.py into WSL.")
156 return None
157 staged_video = self._stage_video(video_win_path, f"{key}{ext}")
158 if staged_video is None:
159 return None
160 frames_dir = f"{self.cfg.wsl_stage_dir}/{key}_frames"
161
162 # Disk guard (D11): a 4K clip can extract ~100 GB of PNG frames. BLOCK before
163 # the WSL stage fs can't hold ~15x the source ? a mid-run disk-fill wastes a
164 # long run and
165 # looks like a quality fail. Best-effort measuring (an undeterminable read never
166 # Streaming never materializes the PNGs, so the guard only matters off the fast
167 # path.
168 if not self.cfg.stream_video:
169 shortfall = self._wsl_frame_stage_shortfall(video_win_path)
170 if shortfall:
171 logger.error("WASB inference aborted ? %s", shortfall)
172 return None
173
174 # Fast path: stream-decode the video in-process (no PNG extraction). Output is
175 # bit-identical to the PNG round-trip but avoids writing ~46k PNGs for a 12-min
176 # clip.
177 if self.cfg.stream_video:
178 frame_args = "--stream-video "
179 else:
180 frame_args = f"--frames_out_dir {wsl_tilde_quote(frames_dir)} "
181 cmd = (
182 f"conda activate {self.cfg.conda_env} && cd {self.cfg.repo_dir}/src && "
183 f"python wasb_infer.py "
184 f"--video {wsl_tilde_quote(staged_video)} "
185 f"{frame_args}"
186 f"--weights {wsl_tilde_quote(self.cfg.weights_path)} "
187 f"--sport {shlex.quote(self.cfg.sport)} "
188 f"--out {wsl_tilde_quote(wsl_out_dir + '/' + key + '_wasb.csv')}}"
189)
190 logger.info("Running WASB inference on %s ...", base)
191 try:
192 res = self._wsl_bash(cmd)
193 except subprocess.TimeoutExpired:
194 logger.error("WASB inference timed out after %ss.", self.cfg.timeout_sec)
195 return None
196 if res.returncode != 0:
197 logger.error(
198 "WASB inference failed:\n%s", (res.stderr or res.stdout)[-2000:]
199)
200 return None
201 if not os.path.exists(expected_csv_win):
202 logger.error(

```

```

200 "WASB finished but expected CSV not found at %s", expected_csv_win
201)
202 return None
203 logger.info("WASB trajectory CSV: %s", expected_csv_win)
204
205 # Success ? the trajectory CSV is the durable output; the staged video copy and
206 # the (huge) decoded frame cache are now pure intermediates. Delete them unless
207 # explicitly retained. On any earlier failure we return above WITHOUT cleaning,
208 # so the frame cache + detector cache remain for a resumable re-run.
209 # (Streaming never writes a frame cache, so there's only the staged video to
drop.)
210 if not self.cfg.keep_frames:
211 stale = (
212 [staged_video] if self.cfg.stream_video else [frames_dir, staged_video]
213)
214 logger.info(
215 "Cache hygiene: removing staged video%s for %s "
216 "(set indexing.wasb.keep_frames=true to retain).",
217 "" if self.cfg.stream_video else " + frame cache",
218 key,
219)
220 self._wsl_rm_rf(stale)
221 return expected_csv_win
222
223 # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
224 parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
225 trajectory_to_action_windows = staticmethod(
226 TrackNetRunner.trajectory_to_action_windows
227)
228

```

## 28. assistant (2026-07-03T19:50:02.361Z)

I now have everything needed for a precise analysis. Let me synthesize the finding.

### Analysis of the tracker-postproc lever

```

The loop (`wasb_infer.py:684-690`, `run()`):
```python
tracker = build_tracker(cfg)
tracker.refresh()
rows = []
for fid in sorted(det_by_fid.keys()):
    r = tracker.update(_dets_to_tracker(det_by_fid[fid]))
    rows.append((fid, 1 if r.get("visi") else 0, r.get("x", -1), r.get("y", -1)))
...

```

****Scale:**** 7.5 min × 60 fps ? 27,000 frames ? ~27,000 windows ? `sorted(det\_by\_fid.keys())` ? 27,000 iterations. `sorted()` is $O(n \log n)$ on ~27k ints ? sub-millisecond, not a bottleneck.

****Is it a serial CPU bottleneck running after the GPU is idle?**** Yes, it is serial-CPU and runs after detection. But the crucial question is its magnitude vs the 727.9s detection pass. The WASB online tracker (`build\_tracker`) is a classic per-frame greedy/Hungarian nearest-neighbor associator over the frame's candidate list. Per `update()`: `\_dets\_to\_tracker` allocates a handful of tiny numpy arrays (one `np.array([x,y])` per candidate ? WASB typically emits 0(1-10) candidates/frame across scales), then the tracker does a few `np.linalg.norm` distance comparisons against its last known position. That's on the order of tens of microseconds per frame. At 27k frames that is well under 1-2 seconds total ? roughly 0.1-0.3% of the 727.9s detection pass. The docstring itself calls this pass "cheap" and "costs seconds, not the GPU hour" (`wasb\_infer.py:28-29`).

****Is it parallelizable/vectorizable?**** No ? and this is the load-bearing point. The tracker is a ****stateful online filter****: `tracker.refresh()` resets state, then each `update(frame\_t)`

depends on the accumulated track state from frames `0..t-1` (last position, velocity, track continuity). This is an inherently sequential recurrence; you cannot split frames across threads/processes without changing the association decisions, and any reordering/vectorization changes the output ? which ****breaks the D5/D16 byte-parity guarantee**** the repo asserts.

****Per-window Python overhead / $O(n^2)$ check:**** I checked the detection-pass assembly too. `frames\_payload` (`wasb\_infer.py:648-657`) is per-batch, $O(\text{batch} \times \text{frames_in})$? linear, no $O(n^2)$. `det\_by\_fid[fid].extend(plain)` is amortized $O(1)$ list append. The JSONL flush (`\_append\_windows`, `wasb\_infer.py:220-225`) writes `pending` and fsyncs every `flush\_every=1000` windows ? $O(n)$ total writes, fsync count = $n/1000$? 27, negligible. `load\_and\_clean\_cache` (only on resume) is a single $O(n)$ linear scan. ****No $O(n^2)$ anywhere.****

Let me verify there isn't a hidden quadratic in how `det\_by\_fid` is consumed vs built, and confirm the tracker output shape once more before finalizing.

29. user (2026-07-03T19:50:06.879Z)

```
118: def _dets_to_tracker(plain_list):
129: def write_manifest(cache_dir: str, manifest: dict) -> None:
193:     Returns (windows_done, det_by_fid) where det_by_fid maps frame_id -> list of
198:     det_by_fid = defaultdict(list)
214:         det_by_fid[int(fid)].extend([list(p) for p in plain])
217:     return len(valid), det_by_fid
220: def _append_windows(fh, objs) -> None:
518:     frames = sorted(
555:     windows_done, det_by_fid = load_and_clean_cache(cache_dir)
563:     windows_done, det_by_fid = 0, defaultdict(list)
576:     write_manifest(cache_dir, base_manifest)
610:     pending = []
649:         frames_payload = []
650:         for ie in sorted(batch_results[ib].keys()):
654:             frames_payload.append([fid, plain])
655:             det_by_fid[fid].extend(plain)
656:             pending.append({"w": win_idx, "f": frames_payload})
659:         if len(pending) >= flush_n or done >= windows_total:
660:             _append_windows(det_f, pending)
661:             pending = []
663:             write_manifest(cache_dir, base_manifest)
674:         if pending:
675:             _append_windows(det_f, pending)
677:             write_manifest(cache_dir, base_manifest)
688:     for fid in sorted(det_by_fid.keys()):
689:         r = tracker.update(_dets_to_tracker(det_by_fid[fid]))
690:         rows.append((fid, 1 if r.get("visi") else 0, r.get("x", -1), r.get("y", -1)))
```

30. assistant (2026-07-03T19:50:14.212Z)

Confirmed. One subtle detail worth noting: `det\_by\_fid[fid]` accumulates candidates across ****every window a frame appears in**** ? because a frame appears in up to `frames\_in` overlapping windows, each frame's candidate list is the concatenation of that frame's detections from all `frames\_in` windows (this is the `test\_replay\_accumulates\_detections\_across\_windows` contract). So `\_dets\_to\_tracker` at frame `fid` rebuilds up to `frames\_in` x `candidates\_per\_window` numpy arrays. `frames\_in` for WASB is small (3 for the badminton model). Still tiny ? a few dozen numpy allocations per frame at most. Total allocations ? $27k \times O(10)$? a few hundred thousand tiny `np.array([x,y])` calls, which is where the only marginal cost lives (Python/numpy allocation overhead, on the order of 1-3s total, not minutes).

My analysis is complete. This lever is ****not**** where the 727.9s is going ? the GPU-idle problem is the CPU-feed serialization during detection (the prefetch/batch lever), not this post-processing tail. Let me record the structured finding.

31. user (2026-07-03T19:50:42.157Z)

Structured output provided successfully