

claude-windows / dbc4cdae-04f6-4d06-af90-397bbce3fc57

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\dbc4cdae-04f6-4d06-af90-397bbce3fc57\subagents\agent-ac396efe913fa3264.jsonl`  
- SHA-256: `224f3ed95dc435acd974513b1b2c9f43b4b98ed5615cd59ddad60c8a8ba893a2`  
- Source modified: `2026-07-06T20:15:48+00:00`  
- Imported at: `2026-07-08T15:59:37+00:00`  
- project: `subagents`  
- session_id: `dbc4cdae-04f6-4d06-af90-397bbce3fc57`
```

Transcript

1. user (2026-07-06T20:13:55.756Z)

Read-only investigation. Do NOT modify any files. I need exact content from several docs in the repo at C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer. For each item, give me the EXACT file path and quote the concrete values (paths, URIs, filenames, resolutions, fps, F1 numbers, commands) with the line region where they appear.

Read these files fully and extract the requested facts:

1. C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\docs\GOLDEN_REGRESSION_FIXTURES.md
 - Where do golden eval VIDEOS live (local paths, GCS gs:// URIs)? List actual filenames/paths.
 - Where are ground-truth LABELS and what format (CSV columns)? How are video<->label paired?
 - Any small subset (2-3 clips) named for a quick run, with resolution/fps.
2. C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\docs\EVALUATION.md
 - The eval command(s) for a single video, and what inputs they need.
 - The F1/precision/recall definition used (rally-level? tolerance? IoU threshold?).
3. C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\docs\DATA_IN_GCS.md
 - GCS bucket names and gs:// URIs for golden videos and labels. List actual bucket/prefix paths.
4. C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\docs\data_pipeline\GOLDEN_VIDEOS.md and GOLDEN_SET_PHASE1_VIDEOS.md and GOLDEN_SET_IMPLEMENTATION.md (all three, in docs\data_pipeline\
 - Actual golden video filenames, GCS paths, resolutions, fps, and the label CSV format (columns).
 - Video<->label pairing convention (e.g. X.MP4 <-> X.rallies.csv).
 - Names of 2-3 specific clips usable for a quick numbers run.
5. C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\docs\R2_EVAL_METRIC_DESIGN.md
 - The exact definition of tolF1 vs F1, and the tau (?) tolerance values (in seconds or pixels?).

Return a concise, structured report organized by the 4 questions: (a) golden video locations/filenames/GCS URIs, (b) ground-truth label locations + CSV format + pairing, (c) 2-3 quick-run clips with resolution/fps, (d) F1/tolF1 definition + ?. Include exact file:line-region evidence. Flag anything not found as an open question. Return findings as text only ? do not write any files.

2. assistant (2026-07-06T20:14:01.491Z)

I'll investigate these documentation files. Let me start by reading all the requested files in parallel.

3. user (2026-07-06T20:14:02.821Z)

```
1      # Golden Regression Fixtures ? video-free, non-attributable regression suite
2
3      > **Status:** `IN PROGRESS` ? owner-gated . **Owner:** Avi . **Created:** 2026-06-16 .
**Last updated:** 2026-06-16 . **Legal:** owner **accepted the researched risk** 2026-06-16
(conditions §3 stand; this is the owner's informed acceptance, **not** a retained-counsel
opinion)
4      > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE, per archive policy)
5      > **Tracking anchors:** §7 progress tracker is the source of truth; indexed in
`docs/README.md`; pointer in memory `current-state` (and `docs/NEXT_STEPS.md` once work starts).
6      > **Relation to existing docs:** EXTENDS (does not supersede)
`docs/GOLDEN_DATA_SHARING.md` (#175); reuses the eval stack documented in `docs/EVALUATION.md` /
`docs/EXPERIMENT_HARNESS.md`; quality floor ties to `docs/QUALITY_ITERATIONS.md`.
7      > **Honesty note:** §3 (Legal) is **researched general information, NOT legal advice** ?
it gates on retained counsel (§8). §2/§4/§7 mark which claims are **[verified]** against code vs
**[design]** proposals.
8
9      ---
10
11     ## 0. TL;DR
12
13     Let a contributor on **another machine with no video, GPU, WSL, or DB** run the rally-
segmentation regression suite from **committed structured files**, so refactors and non-quality
PRs get a red/green side-effect signal. Two layers, built once and grown by accretion:
14
15     1. **Mechanics** ? freeze the deterministic pipeline at the **post-detector trajectory
seam** and replay it through the **already-shipped** eval stack (`experiment.py` / `metrics.py` /
`windowing.py` / `regression.py`), with two auto-discovered gates: a
**characterization/snapshot** gate (catches any behaviour change) and a **quality-floor** gate
(catches quality drops).
16     2. **Privacy/legal filter (SEALED-FIXTURES)** ? only **de-attributed, owner-source,
biometric-free** numbers are ever committed; richer/sensitive streams stay key-gated or out-of-
repo. Non-attributability comes from **deletion-at-source**, not encryption. ? "Minimized" means
**stream minimization** (shuttle-only, no person/audio/pose) ? **NOT** reduction of the shuttle
stream to ~5 scalars: the de-attributed per-frame `trajectory.csv` IS committed (see §4c + its
risk-acceptance re-confirmation note).
17
18     **The convergence that makes this clean:** the **privacy** answer (publish shuttle-only,
drop person/pose streams) and the **determinism** answer (freeze only the `CONTROL` path ? which
runs **no numpy execution**, keeping the numpy/BLAS learned scorer off the cross-machine gate)
point at the **same minimal public tier**. That coincidence is the spine of the design. (numpy
is still **imported** transitively via `experiment` and must be installed; the CONTROL replay just
never calls into it.)
19
20     **Two findings already actionable, independent of the rest:**
21     - ? **Live de-attribution bug** ***(verified)**: real player/venue names are committed
today in `eval_baselines/heuristic_lovo_n6.json`,
`eval_baselines/gemini_wrapper_2026-06-10.json`, and `backend/eval/eval_partial_labels.py` (the
last also has a personal `C:/Users/avidu/...` path) ? and in git history. ? **Deliverable P0.**
22     - ? **Determinism trap** ***(verified, corroborated by two independent red-teams)**: the
learned variants (`shuttle_only`, `fusion`) train a numpy/OpenBLAS `DYNAMIC_ARCH` logistic
regression whose decisions flip near the 0.5 threshold across CPUs ? would false-RED on a
contributor's machine. ? the cross-machine gate must be **CONTROL/static-only**.
23
24     ---
25
26     ## 1. Problem & goal
27
28     The rally-quality measurement is split across two machines by capability
(`docs/GOLDEN_DATA_SHARING.md`):
29     - `backend/eval/fusion_golden.py` **extracts** `_golden_features.json` ? needs the
**video + trajectory CSV + labels + GPU/WSL** (GPU box only).
30     - `backend/eval/fusion_compare.py` / `backend/eval/ablation.py` **re-score** those JSONs
```

```

on **pure CPU, no video** ? any machine.
31
32     The CPU re-scoring is *already* video-free, but the inputs live in gitignored `output/`
/ private OneDrive, so the regression tests that depend on them **skip everywhere**
(`tests/test_ablation.py::test_litmus_reproduces_gate_a`, `tests/test_fusion_compare.py`). The
committed `eval_baselines/` baselines exist but `regression.py` still scores DB-stored
predictions (needs a pipeline run on the video).
33
34     **Goal:** commit a tiny, version-tagged, **numeric-only, non-attributable** per-fixture
corpus + auto-discovered gates so a clean `git clone` / CI gets a real red/green regression
signal with **zero video/GPU/DB**, while never leaking attributable or sensitive data.
35
36     ---
37
38     ## 2. Decisions locked
39
40     | # | Decision | Source | Implication |
41     |---|-----|-----|-----|
42     | D1 | **Corpus is owner-recorded / owned** | owner answer 2026-06-16 | Legal analysis
takes **Fork A** (clean): no third-party copyright/ToS infringement; concern is trade-
secret/moat + privacy. Strongest non-attributability (no public reference to correlate against).
|
43     | D2 | **Repo private now, may open-source later** | owner answer 2026-06-16 | Design
for **public-grade hygiene from day one** ? git history is permanent; nothing scrubable-later. |
44     | D3 | **Adults, consent not formalized** | owner answer 2026-06-16 | **Person/pose/gait
streams stay OUT of the public tier** (DPDP/GDPR derived-personal-data risk without consent).
Public tier = **shuttle + abstract intervals + non-biometric features only**. |
45     | D4 | **Freeze seam = post-detector trajectory** *(design; SHIPPED differs ? see ?)* |
workflow-1 synthesis | Widest deterministic surface; reuse the shipped eval stack. ?
**Correction (audit 2026-06-16):** the shipped freeze **commits the de-attributed per-frame
(Frame,Visibility,X,Y)` `trajectory.csv`** (shuttle-only, time-origin reset) ? it is **NOT**
minimized to ~5 aggregate scalars. See §4c (and its risk-acceptance re-confirmation note). |
46     | D5 | **Cross-machine gate is CONTROL/static-only** *(design)* | both red-teams
[verified cause] | Learned variants (numpy/BLAS) are **owner-box/Tier-1 only**, asserted with
tolerance, never as a cross-machine byte-snapshot. |
47     | D6 | **Non-attributability = deletion-at-source, encryption = public-only barrier**
*(design)* | workflow-2 synthesis | A contributor who runs the tests necessarily reads
plaintext+key (DRM/analog-hole). So we **delete** attribution, not hide it. Encryption only
stops the no-key public. |
48     | D7 | **Two tiers** *(design)* | workflow-2 synthesis | **Tier-0** public/no-key
(shuttle-only, de-attributed). **Tier-1** key-gated/out-of-repo (person/audio, full corpus),
**skips cleanly when absent**. |
49
50     ---
51
52     ## 3. Legal foundation (researched ? gates on counsel, §8)
53
54     **Core verdict (settled across IN/US/EU):** the derived numbers ? per-frame shuttle
(x,y), rally start/end intervals, audio/court/aggregate-person feature values ? **do NOT inherit
the source video's copyright.** They are uncopyrightable *facts/measurements*, not authored
expression, and not a §101 "derivative work."
55
56     | Jurisdiction | Copyright of the data | Database right | Contract / ToS | Net |
57     |---|---|---|---|---|
58     | **India** (primary) | Not copyrightable as facts (*EBC v. D.B. Modak*); but
compilations get thin protection on a low "skill/labour" bar (*Burlington*, *OLX v. Padawan*) ?
only bites if **lifted from a 3rd-party dataset**, not self-derived | **None** (no sui-generis
right) ? a real permissiveness edge | Anti-scraping ToS (untested in court) + **IT Act
s.43/s.66** civil/criminal exposure for 3rd-party ingestion | **Owned-source + minors-excluded +
no person/pose = clean** |
59     | **US** | Not copyrightable (*Feist*; §102(b); *NBA v. Motorola*) | None (*Feist*
killed sweat-of-the-brow) | **Dominant risk.** ToS enforceable where copyright/CFAA fail
(*ProCD*, *hiQ* ? injunction + **forced deletion**). Remedy for redistribution is takedown, not
nominal damages | Publishable as facts **iff** owned/permissioned source + no barring ToS + no
person/biometric/minor stream |

```

60 | ****EU / UK**** | Not copyrightable (**Football Dataco v. Yahoo**) | ****Genuinely**
 contestable.**** *Football Dataco v. Sportradar*:** recording pre-existing on-court facts =
"obtaining**"**, not "creating" ? our "we created it" premise is ****overconfident****; harm test =
**CV-Online Latvia* | Ryanair v. PR Aviation: ToS can ban extraction/redistribution even with no*
IP right | Get EU+UK advice before redistributing numeric streams at scale |
 61

62 ****Provenance is the decisive fork.**** ****Fork A (owner-recorded, unpublished)**** ? the only
 path safe to commit publicly; collapses to a trade-secret **choice** (Indian breach-of-confidence
 / Contract Act 1872 / IT Act s.72 ? ****not**** US DTSA) + privacy. ****Fork B**
 (broadcast/YouTube/scraped)****** ? do ****not**** commit; the extraction act needs a license or
 now-****contested**** US fair use (**Thomson Reuters v. Ross* 2025; USCO May-2025*), and unlawful
 acquisition independently sinks it (**Bartz v. Anthropic**, ~\$1.5B settlement). Caveat: owner
 footage ****shot at a third party's event**** (tournament/club) carries
 ticket/venue/organisier/broadcaster terms ? treat as Fork B until cleared.
 63

64 ****Publish-by-stream (privacy):****
 65 - ? ****Publish-OK:**** shuttle (x,y)+Visibility (inanimate object), abstract rally/GT
****intervals**** (identity-stripped), non-biometric audio/court/aggregate-person features detached
 from identity.
 66 - ?? ****Conditional (effectively never in v1):**** person bounding boxes that **persistently*
 single out a player ? flips to GDPR Art.9 the moment it identifies. This architecture is prone
 to that flip ? treat per-player tracking as a tripwire.
 67 - ? ****Never publish:**** pose/skeleton/****gait**** (re-identifies at 80?95% even with faces
 blurred ? **Weiss et al. 2025**), face crops, any named-player?movement mapping. Matches the
 existing "gait/biometrics strictly future (GDPR Art.9)" guardrail.
 68 - ? ****Minors flip everything:**** DPDP "child" = under 18; s.9(3) ****prohibits behavioural**
monitoring** (the Rule-12 exemptions don't cover this project). A persistent per-player
 signature on a minor is, conservatively, prohibited. ****Exclude minors**** from any persisted
 stream beyond inanimate shuttle/interval data.
 69

70 ****Non-attributability is relative, not absolute.**** A per-video trajectory is a
 fingerprint (de Montjoye: 4 points ? 95% re-ID). It's only non-attributable because the ****Fork-A**
 source stays unpublished (no public reference to correlate against). Per **EDPS v. SRB** (CJEU
 C-413/23 P, 2025) it ****remains personal data to the holder**** who can re-run the detector. So:
 conditional, revocable (publish the source and the link snaps back), forward-looking (re-assess
 as the public-video environment grows). Never market it as "anonymous."
 71

72 ***Key authorities:** Feist 499 U.S. 340; NBA v. Motorola 105 F.3d 841; 17 U.S.C.
 §101/§102(b); Sega v. Accolade 977 F.2d 1510; Football Dataco v. Yahoo (C-604/10) & v.
 Sportradar (C-173/11); BHB v. William Hill (C-203/02); CV-Online Latvia (C-762/19); Ryanair v.
 PR Aviation (C-30/14); Thomson Reuters v. Ross (D. Del. 2025); EBC v. D.B. Modak; DPDP Act 2023
 s.3/s.9 + Rules 2025; EDPS v. SRB (C-413/23 P). Full URLs in the research artifact (§10).
 73

74 ---
 75

76 **## 4. Architecture (design)**
 77

78 **### 4a. Freeze seam + reuse map**
 79 Freeze at the ****trajectory CSV ? CONTROL-path replay****. One shared helper
`backend/eval/golden\_fixtures.py::replay\_fixture()` is the **single** code path used by both the
 freeze tool and both gates. As ****shipped****, that path is (per the module docstring):
 80 *`parse_trajectory_csv ? distill_local.build_candidates(proposal=(vt,mg,mwd)) ?*
VideoCandidates(5 shuttle cols only) ? experiment.predict(CONTROL) [keep all ? merge_overlaps] ?
metrics.score`.
 81 Reuses ****verbatim****: *`metrics.score/interval\_iou`,*
`experiment.predict/CONTROL/VideoCandidates`, `distill\_local.build\_candidates`,
`tracknet\_runner.parse\_trajectory\_csv`/`trajectory\_to\_action\_windows`, `rally\_gate` net-
crossings (inside `build\_candidates`), `gt\_loader.load\_gt\_intervals` (strict),
`regression.gt\_hash`, `eval\_baselines/` committed-baseline precedent.
 82 > ? ****Correction (audit 2026-06-16):**** the replay does ****NOT**** import *`windowing.py`* ?
 there is no *`windowing.build\_windows(preset)`* / *`WindowingPreset`* / "default SERVED" call. The
 windowing params travel as a ****raw 3-tuple `(vt, mg, mwd)`**** pinned in each fixture's
`provenance.json` and passed straight to *`build\_candidates`*; the harness deliberately avoids the
 named *`windowing.SERVED`* constant so a future SERVED retune can't silently churn fixtures.
 (*`segmentation\_metrics`, `experiment.compare/Variant.spec\_hash`, `regression.save\_baseline`,*

```

`eval_stats` are not on the M1 replay path either ? the quality-floor gate M4 will add the
baseline/stats pieces.)
83
84     ### 4b. Tiered model (SEALED-FIXTURES)
85     | Tier | Where | Key? | Contents | Test behaviour |
86     |---|---|---|---|---|
87     | **Tier-0** | `eval_baselines/fixtures/` (tracked) | none | opaque surrogate slug,
fps/frame_width (retained, NOT quantized), the de-attributed per-frame `trajectory.csv`
(Frame,Visibility,X,Y) (shuttle-only, time-origin reset), the 5 shuttle feature columns inside
`expected.json`, strict-slug `label`, relative `gts`. **No person/audio/pose.** ? NOT minimized
to 5 scalars ? see §4c. | **Unconditional** ? runs on every PR/fork; turns today's skipped
litmus into a real public gate |
88     | **Tier-1** | OneDrive (default) **or** SOPS-encrypted in-repo | token/age key | full
corpus **with** person/audio; authoritative ? F1/CI/sign-test + exact Gate-A litmus | **Skips
cleanly **when key/data absent (generalize `_golden_present()` ? `_tier1_present()` ) |
89
90     A fresh no-key clone is always **green**; coverage never depends on a secret.
91
92     ### 4c. Anti-attribution measures (Tier-0)
93
94     > ? **RISK-ACCEPTANCE ? OWNER RE-CONFIRMED option (a), 2026-06-22.** The shipped freeze
(`freeze_real_fixture`) commits the de-attributed per-frame `(Frame,Visibility,X,Y)`
`trajectory.csv` (time-origin reset + opaque slug + metadata strip + scenario-only label) ?
**NOT** minimized to ~5 scalars. A per-frame stream is a de Montjoye fingerprint, so non-
attributability rests on the **source staying unpublished** (EDPS v. SRB), not on minimization.
(An earlier draft wrongly claimed "5 scalars / never the per-frame CSV"; the audit corrected it
2026-06-16.) The owner cleared ALL owned golden videos (owned + minors-free) and accepted the
corrected per-frame basis on 2026-06-22 ? the **6 real fixtures shipped under P3**. Owner
alternative ? true ~5-scalar minimization ? remains available if these ever go public.
95
96     - **Opaque RANDOM surrogate id** (`os.urandom` + a private surrogate?source map kept
off-git) ? **not** the MD5 `video_id`, and **not** HMAC(salt, video_id?name) (red-team: brute-
forceable over a tiny known roster if the salt leaks). Never persist the source `name`/filename.
97     - **Strict-slug `label` + dropped `source_note`** ? the operator `label` is validated
against `^[a-z0-9][a-z0-9\-\_]*$` (a free-text identifier like a player/event name is refused),
and operator free-text `source_note` is **accepted-but-NOT-persisted** (never reaches the
committed `provenance.json`) ? both close the free-text identifier leak `_scan_forbidden` could
not police (audit fix #1).
98     - **Metadata strip** ? no filename, video_id, match/event/player identity, capture date.
99     - **Time-origin reset** ? subtract per-fixture `t0` (6dp-rounded) from all
start/end/gts. **Metric-invariant within `_FLOAT_TOL`** (`metrics.interval_iou` and start/end-
error are pure differences): the genuine (raw) reset-vs-no-reset score delta sits at the float-
error floor (~3.3e-7); bit-identical for whole-second offsets. See §4c-note below and §6.
100    - **Per-frame trajectory IS committed (de-attributed), NOT minimized to scalars** ?
Tier-0 commits the de-attributed per-frame `trajectory.csv` (shuttle stream only). It is **not**
reduced to ~5 aggregate scalars; the 5 shuttle **feature** columns appear (computed) inside
`expected.json`, but the raw per-frame `(Frame,Visibility,X,Y)` stream is part of the committed
fixture. (This is the corrected basis the ? note above gates on.)
101    - **Stream minimization (streams, not resolution)** ? Tier-0 keeps the shuttle block
only; person/audio excluded; pose/gait/face have **no field in the schema by design** (positive
`_scan_forbidden` assertion on the written bytes).
102    - **fps/frame_width are RETAINED as scale constants** ? they are NOT quantized today
(quantization is a deferred P3 option, see §7 ?). Reconcile any "quantize fps/frame_width"
wording elsewhere to this.
103    - **Provenance hard-fail gate** ? the de-attribution pass refuses to emit a public
fixture for any record not tagged owner-recorded(Fork-A) + minors-excluded + non-blank-license,
and refuses a negative reset label (timeline predates the trajectory). Writes go to a temp dir
and are promoted only after the privacy scan passes, so a refused freeze leaves no partial
fixture (audit fixes #3a/#3b).
104
105    ### 4d. The two gates
106    - **Characterization/snapshot** ? replay fixture, **parse-compare** JSON (never byte-
compare ? CRLF/float-format safe; `core.autocrlf=true` here), exact on int/structural fields,
`approx(abs=1e-6)` on floats; stamp + check windowing-preset/label/gt fingerprints
(incommensurable ? loud fail). Catches any behaviour change even at constant F1. **CONTROL path

```

```

only** across machines.
107     - **Quality-floor** ? score vs committed labels; assert per-video + mean F1 ? floor in
`eval_baselines/regression_baseline.json` (created here; carries a **windowing fingerprint** so
a served-default change is flagged incommensurable, not silently compared); emit `eval_stats`
paired-delta-CI + sign test (NUMBERS INFORM); add a **metrics-vs-base-branch** paired delta.
Synthetic-tier floors measure **correctness, not field quality** (documented).
108
109     ### 4e. Serialization & optional protobuf
110     JSON with `sort_keys=True` + pre-quantized floats + LF is sufficient. **Protobuf is
optional (01)** ? adopt only if cross-OS byte-stable snapshots are wanted (proto3
`rally.fixture.v1`, codegen + `protoc && git diff --exit-code` guard, `*.pb binary -text`).
Protobuf is **encoding, not secrecy** (trivially `--decode_raw`).
111
112     ### 4f. Encryption & keys (Tier-1 only)
113     **Default = out-of-repo + token** (existing OneDrive `RALLY_GOLDEN_DIR` seam ? zero
sensitive bytes in git history, sidesteps the legal "redistribution" act). **Alternative = SOPS
+ age** (per-recipient, revocable, value-level diffs) for Fork-A-clean-but-private data wanting
in-repo versioning. Rejected: git-crypt (whole-repo re-key), git-lfs (storage ? secrecy). CI
secret exposed **only to the Tier-1 job on trusted branches, never to fork-PR runners**. The
**HMAC/surrogate salt** (if used) and the surrogate?source map live with the private corpus,
rotated per release. Document explicitly that Tier-1 crypto defends **only** the no-key public.
114
115     ---
116
117     ## 5. Threat model (who's protected vs not)
118
119     | Actor | Protected against | NOT protected (by design / residual) |
120     |---|---|---|
121     | **TM1 no-key public** | map fixture?source/match/player; decrypt Tier-1; confirm a
source even holding the candidate file | read Tier-0 numbers + run shuttle-only harness
(intended; useless-in-isolation) |
122     | **TM2 keyed contributor** | recover filename/video_id/match/face/pose/gait ? *never
written* (survives the analog hole) | read every value the suite reads (unpreventable; not
pretended otherwise) |
123     | **TM3 CI runner** | leaked log can't leak attribution it never had | runs the suite
(intended); Tier-1 secret never on fork runners |
124     | **TM4 malicious keyed insider** | correlation needs a public reference; Fork-A source
is unpublished ? "not reasonably likely" | **residual:** conditional + revocable ? publish/leak
the source later and the link can snap back (forward-looking re-test duty) |
125     | **TM5 re-ID adversary** | source unpublished + reset timeline ? no back-match against
a public reference | ? becomes live if a Fork-A source is ever published ? **and the committed
per-frame trajectory IS a fingerprint** (de Montjoye 4-point re-ID), NOT just ~5 scalars
(corrected S4c). The protection is the unpublished source, not minimization |
126     | **TM6 accidental plaintext commit** | required CI leak-scan rejects MD5/known-
name/non-surrogate/plaintext-SOPS | if the required check is disabled, a leak lands |
127
128     ---
129
130     ## 6. Honest limits ? what a green suite does NOT prove
131
132     - Green proves only that the **deterministic post-trajectory transforms** are unchanged
for the **frozen cases**. It does **not** cover: the WASB/TrackNet detector,
`_probe_fps_width`/cv2 frame-seek, the AI/provider handoff, chunking/re-encode, the ffmpeg
stitch, DB persistence/migrations, or the API filter. *A refactor that breaks decoding or the
detector passes these gates green.* The owner-box `regression.py` DB run remains the only full-
chain check.
133     - Characterization = **behaviour-locking, not correctness** ? it freezes current
behaviour *including current bugs*. "Passing" = "behaviour unchanged", never "behaviour
correct". This caveat must ride in the **test pass/fail message**, not only the docs.
134     - **Synthetic-tier floors = plumbing correctness, NOT field quality.** Never cite
synthetic F1 in `QUALITY_ITERATIONS.md` / the paper as rally-detection performance (tag `kind:
synthetic`).
135     - Tier-0 public coverage is **shuttle-only** ; the person/audio fusion ?F1 and exact
Gate-A values need Tier-1 (skipped on no-key clones).
136     - Data is **not "anonymous" absolutely (EDPS v. SRB) ? relative, conditional,
```

revocable, forward-looking.

137 - The committed Tier-0 stream is the **de-attributed** per-frame `trajectory.csv`, NOT ~5 scalars ? a per-frame trajectory is a de Montjoye fingerprint; the only protection is the **unpublished source** (§4c ?, §5 TM5).

138 - **Time-origin** reset is metric-invariant within `_FLOAT_TOL`, not byte-exact for every offset. The genuine (raw) reset-vs-no-reset score delta is at the float-error floor (~3.3e-7); it is **bit-identical** for whole-second offsets, but because 6dp-rounding is not additive (`round6(a)-round6(b) ? round6(a-b)`) the **serialized** 6dp boundary error can differ by one 6dp ULP (1e-6) on a sub-second offset ? the serialization grid, not a metric divergence. (Audit fix #2 rounds `t0`+labels to 6dp before the shift to keep the residual ? the grid step.)

139 - Encryption gives **no secrecy** from a running contributor; protobuf gives **no secrecy** at all.

140 - De-attribution does **not** cure provenance/licensing ? it sits **behind** the per-record provenance gate + counsel sign-off.

141 - `--refresh-snapshot` can rubber-stamp a real regression; mitigation is **review** of the visible diff + a separate, deliberate floor-baseline update. **A recorded --reason** is **REQUIRED** (the CLI errors without one ? audit fix #6).

142

143 ---

144

145 ## 7. Deliverables & progress tracker ? **source of truth**

146

147 Legend: ? Todo · ? In progress · ? Done · ? Blocked/gated. One **small PR per row**, branched from `origin/master`, no stacking. Update the row (status + PR link) as work lands.

148

ID	Deliverable	Depends on	Counsel-gated?	Status	PR
M1	Fixture framework + synthetic Tier-0 + shared replay helper (also delivered M2-core + M3 ? see below). <code>backend/eval/golden_fixtures.py</code> (<code>replay_fixture</code> , schema dispatch, <code>Path(__file__)</code> -relative) + 3 tool-generated synthetic fixtures + manifest + README + <code>.gitattributes</code> LF pin + <code>tests/test_golden_fixtures.py</code> . CONTROL-only (no numpy execution on the replay path); cross-OS verified. ? No (synthetic) ?				
M2	Freeze/refresh extractor. Core (<code>freeze_synthetic</code> / <code>refresh_snapshot</code> / <code>--check</code> CLI + idempotence test) shipped in M1 (#189); the <code>--license</code> / <code>--minors-free</code> / <code>--owned</code> refusal gate shipped with P1 (this PR). M1 No ?				
M3	Characterization gate ? <code>tests/test_golden_fixtures.py</code> : parse-compared (exact-int / approx-float 1e-6), CONTROL-only , positive no-person/audio assertion, perturbation + idempotence. Shipped in M1 (#189) . (Windowing/label fingerprint guard + row-count cap ? fold into M4.) M1 No ?				
MX	Cross-OS determinism guard (CI). Lightweight <code>golden-crossos</code> CI job (ubuntu/windows/macOS) running <code>--check</code> + the fixtures test ? locks the cross-OS guarantee (CONTROL replay is numpy-only, no GPU stack; a separate fast job , not full-suite×3). M1 No ?				
M4	Quality-floor gate + first <code>regression_baseline.json</code> . <code>tests/test_golden_quality_floor.py</code> ; create <code>eval_baselines/regression_baseline.json</code> with windowing fingerprint ; metrics-vs-base paired-delta; corpus-level (not per-slug) LOVO for runtime. M1, M2 No ? ?				
P1	De-attribution + minimization ? <code>golden_fixtures.freeze_real_fixture()</code> (extends the M1 module, not a separate script): opaque random <code>fx_hex</code> surrogate slug, metadata strip, metric-invariant time-origin reset , strict GT loader, positive no-person/audio/MD5/source-path assertion , manifest non-clobber; <code>--freeze-real</code> CLI with the refusal gate (exit 2 unless owned + minors-free + license). Tooling-only, 12 synthetic tests; no real data committed . M1 No (tooling) ?				

149 | ID | Deliverable | Depends on | Counsel-gated? | Status | PR |

150 |----|-----|-----|-----|-----|----|

151 | **M1** | **Fixture framework + synthetic Tier-0 + shared replay helper** (also delivered M2-core + M3 ? see below). `backend/eval/golden_fixtures.py` (`replay_fixture`, schema dispatch, `Path(__file__)`-relative) + 3 **tool-generated** synthetic fixtures + manifest + README + `.gitattributes` LF pin + `tests/test_golden_fixtures.py`. **CONTROL-only** (no numpy execution on the replay path); cross-OS verified. | ? | No (synthetic) | ? |

[#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) |

153 | **M2** | **Freeze/refresh extractor**. Core (`freeze_synthetic` / `refresh_snapshot` / `--check` CLI + idempotence test) shipped in M1 (#189); the `--license` / `--minors-free` / `--owned` **refusal gate** shipped with P1 (this PR). | M1 | No | ? |

[#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) · `_this PR` |

154 | **M3** | **Characterization gate** ? `tests/test_golden_fixtures.py`: parse-compared (exact-int / approx-float 1e-6), **CONTROL-only**, positive no-person/audio assertion, perturbation + idempotence. **Shipped in M1 (#189)**. (Windowing/label fingerprint guard + row-count cap ? fold into M4.) | M1 | No | ? |

[#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) |

155 | **MX** | **Cross-OS determinism guard (CI)**. **Lightweight** `golden-crossos` CI job (ubuntu/windows/macOS) running `--check` + the fixtures test ? locks the cross-OS guarantee (CONTROL replay is numpy-only, no GPU stack; a **separate fast job**, not full-suite×3). | M1 | No | ? |

[#191](https://github.com/avidullu/badminton-highlight-indexer/pull/191) |

156 | **M4** | **Quality-floor gate + first** `regression_baseline.json`. `tests/test_golden_quality_floor.py`; create `eval_baselines/regression_baseline.json` with **windowing fingerprint**; metrics-vs-base paired-delta; corpus-level (not per-slug) LOVO for runtime. | M1, M2 | No | ? | ? |

157 | **P1** | **De-attribution + minimization** ? `golden_fixtures.freeze_real_fixture()` (extends the M1 module, not a separate script): **opaque random** `fx_hex` surrogate slug, metadata strip, metric-invariant **time-origin reset**, strict GT loader, **positive no-person/audio/MD5/source-path assertion**, manifest non-clobber; `--freeze-real` CLI with the **refusal gate** (exit 2 unless owned + minors-free + license). Tooling-only, 12 synthetic tests; **no real data committed**. | M1 | No (tooling) | ? | `_this PR` |

158 | ****P2**** | ****Owner-accepted (researched-risk), 2026-06-16.**** Owner accepted the researched IP/privacy analysis (§3) to proceed ****under its conditions**** (owned/Fork-A source · minors excluded · biometric/person streams excluded · de-attributed). Owner's informed risk acceptance ? ****NOT a retained-counsel opinion****; optional future counsel confirmation per §8. | P1 | Owner-accepted | ? | ? |

159 | ****P3**** | ****Tier-0 REAL subset ? SHIPPED.**** Committed ****6** curated, trimmed (?12k-frame ? few-min), de-attributed** real fixtures (`fx\_r01`?`fx\_r06`: single-court / multi-court×2 / doubles×2 / continuously-active blob; both 29.97 & 59.94 fps) via `freeze\_real\_fixture --max-frames` (opaque slug, time-origin reset, scenario-only labels, ****NO**** source name/path/video_id ? grep-verified + leak-scan clean). ~1.9 MB total; the characterization + cross-OS gates now run on real data too. Owner cleared ALL owned golden videos (owned + minors-free) + re-confirmed option (a) (2026-06-22). The surrogate?source map stays ****off-git**** (owner-held). | P1, P2 | conditions §3 ? | ? | _this PR_ |

160 | ****O1**** | ***(optional)* **Protobuf + deterministic serialization.**** `rally.fixture.v1` proto3, codegen + `protoc`/`git diff --exit-code` guard, canonical writer, `.gitattributes` binary. Only if cross-OS byte-stable snapshots are wanted. | M1 | No | ? | ? |

161 | ****O2**** | ***(optional)* **Tier-1 key-gated path.**** `scripts/fetch\_golden.py` (OneDrive token, default) and/or SOPS+age; `\_tier1\_present()`; `RALLY\_GOLDEN\_DIR` wired; Tier-1 CI job gated on secret, never on fork runners. | M1 | No | ? | ? |

162 | ****DOC**** | ****Discoverability cross-links**** ? added the `eval\_baselines/fixtures/` row to `CODE\_MAP §0` + a `golden\_fixtures.py` entry to `CODE\_MAP §2.3`; cross-linked `GOLDEN\_DATA\_SHARING.md` (committed subset ? private OneDrive corpus). README indexed in #186; memory `current-state` kept live. (EVALUATION/QUALITY_ITERATIONS cross-links ****deferred**** ? that doc surface is hot in the concurrent quality track; avoid collision.) | M1 | No | ? | _this PR_ |

163

164 ****Must-fix-before-any-public-commit (red-team), folded into the rows above:**** ? numeric-tolerance not byte-equality + multi-OS matrix + CONTROL-only (M3/D5/MX ?); ? scrub names in ****all**** files + history decision (P0); ? leak-scan a ****required**** check, hooks are convenience only (P0); ? person-optional loader + ****positive**** no-person assertion, don't trust the silent `0.0` default (P1 ?); ? random surrogate id not HMAC-over-roster (P1 ?); ? quantize/drop fps/frame_width (P1 ? fps/fw retained as scale constants; quantize is a P3 option); ? document Tier-0 omits the fusion path (DOC); ? counsel + Fork-A + minors/biometric exclusion as hard preconditions (P2 owner-accepted; enforced in code by P1's refusal gate ?).

165

166 **### Progress log (quantitative ? coverage must hold/improve, CI incl. the 3-OS cross-OS job green)**

167 | PR | Deliverable(s) | Suite | Coverage | Notes |

168 |----|-----|-----|-----|-----|

169 | [#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) | M1 + M2-core + M3 | 746 pass / 7 skip | 80.41% | synthetic Tier-0 + CONTROL replay + characterization; cross-OS proven (CI Linux replayed Windows-frozen fixtures) |

170 | [#191](https://github.com/avidullu/badminton-highlight-indexer/pull/191) | MX + tracker/legal reconcile | full suite green | 80.41% | `golden-crossos` CI job green on ubuntu+windows+macos ? cross-OS ****enforced****; legal ? owner-accepted |

171 | _this PR_ | P1 + M2 refusal gate | 759 pass / 7 skip | ****80.49%**** (+0.08) | `freeze\_real\_fixture` + refusal gate; ****+13 tests****; tooling-only, no real data committed |

172 | _this PR (audit-hardening)_ | adversarial-audit fixes #17#17 | 784 pass / 7 skip | ****80.71%**** (+0.22) | strict-slug `label` + dropped `source\_note` (#1); 6dp-rounded reset within `\_FLOAT\_TOL` (#2); temp-dir-then-promote + negative-label refusal (#3); exact-membership source-path scan (#4); key-set-equality in `compare\_expected` (#5); `--reason` mandatory (#6); independent schema guard (#7); `max\_win` recorded-not-applied note (#8); numpy-execution wording (#9); +failure-semantics & by-kind tests (#10/#11); `test\_golden\_real\_fixtures.py` added to `golden-crossos` CI (#12); doc reconciled to shipped reality incl. per-frame-CSV honesty + risk re-confirm (#13?#17). Synthetic `--check` byte-identical (unchanged). |

173 | _this PR (P0 safe scope)_ | personal-path scrub + leak-scan guard | 1365 pass / 7 skip | 85.27% | scrubbed `C:\Users\?` paths from 3 code docstrings (`eval\_partial\_labels.py` / `gemini\_refine.py` / `README.md`); `tools/leak\_scan.py` (32-hex MD5 + personal-path structural guard + gitignored `.leakscan\_private` name-list) + a CI `leak-scan` job + 10 tests. Functional video-name rename DEFERRED (owner safe-scope). |

174 | _this PR (P3 real fixtures)_ | 6 curated trimmed de-attributed real Tier-0 fixtures + `--max-frames` trim | 1390 pass / 7 skip | 85.27% | `fx\_r01`?`fx\_r06` (~1.9 MB; single/multi-court/doubles/blob, both fps) frozen via `freeze\_real\_fixture --max-frames=12000`; de-attribution grep-verified + leak-scan clean; +trim test; owner re-confirmed option (a). F1 0.0?0.59 = real CONTROL-path (behaviour-locking, ****NOT**** field quality). |


```

175
176 ---
177
178 ## 8. Open owner / counsel questions
179
180 **Counsel questions ? documented basis (the owner accepted the researched risk on
2026-06-16; see P2). These remain the record, and a retained-counsel confirmation is
optional/recommended before scaling beyond a small owned, minors-free, de-attributed subset:**
181 1. Does the EU/UK **sui-generis database right** subsist in our self-derived
shuttle/rally streams on the *Sportradar* "obtained vs created" reasoning, and does
redistributing a de-attributed Tier-0 subset cause *CV-Online* harm?
182 2. Post *Thomson Reuters v. Ross* + USCO-2025, is **US fair use for the extraction
step** defensible if the dataset could compete in a "data/analytics" market ? does publishing
only numbers (never footage) matter?
183 3. India: are anti-scraping ToS enforceable; **IT Act s.43/s.66** exposure for any 3rd-
party ingestion; does pending *ANI v. OpenAI* (~Apr 2026) move it? Confirm **DPDP s.9/Rule 10**
and whether per-player tracking is prohibited s.9(3) monitoring of a minor.
184 4. Sign-off on the **final committed Tier-0 subset + provenance log** ; is non-
attributability-via-private-source a defensible posture to assert given its
conditional/revocable nature?
185
186 **Owner (strategic):**
187 5. Publishing the Fork-A subset **forfeits the trade-secret moat** (Indian breach-of-
confidence framing) ? weighed against community/benchmark/paper aims. Your call.
188 6. **Forward-looking commitment:** do not publish a highlight reel of the *same* source
whose fixtures are committed (without re-clearing), and accept a periodic re-assessment as the
public-video environment grows.
189 7. **Scope of first public commit:** synthetic-only (zero provenance risk, recommended)
vs a counsel-cleared real subset.
190 8. **Tier-1 mechanism:** out-of-repo+token (recommended) vs SOPS+age in-repo. And: is
protobuf (01) wanted, or is canonical-JSON enough?
191
192 ---
193
194 ## 9. Definition of done
195
196 The project is **DONE** (flip `Status ? DONE`, archive) when:
197 - [ ] **P0, M1?M4, P1, DOC** are ? and merged.
198 - [ ] A clean `git clone` on **win + ubuntu + mac** runs `python run_all_tests.py` and
the **Tier-0 gates pass green with no key / no video / no GPU / no DB**.
199 - [ ] The **CI leak-scan is a required status check** and passes (no MD5 / real-name /
non-surrogate / plaintext-SOPS in the tree).
200 - [ ] Characterization + quality-floor gates are **parse-compared, CONTROL-only,
windowing-fingerprinted**, and demonstrably catch a seeded behaviour change.
201 - [ ] `docs/CODE_MAP.md $0` + TEST MAP updated; this doc cross-linked; `NEXT_STEPS.md` +
memory `current-state` point here.
202 - [ ] **Either** the public commit is synthetic-only, **or** (**P2 owner-accepted ? +
P3**) for a real-data commit, under the $3 conditions enforced **in code** by P1's refusal gate.
203 - [ ] $6 honest-limits caveat rides in the test pass/fail messages, not only the docs.
204
205 *Optional 01/02 may remain ? without blocking DONE (they're earned, not assumed).*
206
207 ---
208
209 ## 10. References
210
211 **Internal code anchors (verified):** `backend/eval/{regression.py,fusion_golden.py,fusi
on_compare.py,ablation.py,experiment.py,metrics.py>windowing.py,distill_local.py,gt_loader.py,ev
al_stats.py}` · `backend/pipeline/detectors/{base.py,tracknet_runner.py,rally_gate.py}` ·
`backend/utils/hashing.py::compute_video_id` · `eval_baselines/{heuristic_lovo_n6.json,gemini_wr
apper_2026-06-10.json,experiment_leaderboard.json}` ·
`tests/test_ablation.py::test_litmus_reproduces_gate_a` · `.github/workflows/ci.yml`.
212 **Internal docs:** `GOLDEN_DATA_SHARING.md` (#175), `EVALUATION.md`,
`EXPERIMENT_HARNESS.md`, `QUALITY_ITERATIONS.md`, `MULTI_SIGNAL_FUSION_PLAN.md`, `CODE_MAP.md
$0`.

```

213 **Legal authorities:** see §3 (full URLs in the research workflow artifact). *This
section is researched general information, not legal advice.*
214
215 ---
216
217 ### Changelog
218 - **2026-06-16** ? Initial draft synthesized from two multi-agent workflows (mechanics
design + 2 critiques; legal memo IN/US/EU + SEALED-FIXTURES architecture + red-team). Verified
the live name-leak and the BLAS-determinism cause against code. Status: DRAFT, owner-gated.
219

4. assistant (2026-07-06T20:14:04.638Z)

Let me read the remaining files in parallel.

5. user (2026-07-06T20:14:05.679Z)

```
1      # Annotation & Evaluation
2
3      How to measure how well the pipeline detects rallies, using hand-annotated
4      ground-truth timestamps. The harness scores predictions **already stored in the
5      SQLite DB** ? it makes no API calls and is fully deterministic, so you can run it
6      as often as you like for free.
7
8      ---
9
10     ## 1. What you annotate (and what it does / doesn't do)
11
12     You annotate, per video, the `[start, end)` time range of each **rally**. That's
13     all the evaluation harness needs.
14
15     > **Important:** these rally timestamps measure and tune **rally segmentation**
16     > ? i.e. *where* the pipeline says rallies begin and end. They do **not** train
17     > the shuttle detector (TrackNet/WASB), which learns from per-frame shuttle
18     > `(x, y)` coordinate labels, a different and much heavier annotation. Use rally
19     > timestamps as your **golden evaluation set**; only invest in frame-level
20     > coordinate labels if evaluation shows *detection* (not segmentation) is the
21     > bottleneck. See `docs/TRACKNET_WSL_SETUP.md`.
22
23     ## 2. Annotation file format (CSV)
24
25     One CSV per video, one rally per row:
26
27     ```csv
28     start,end
29     00:01:12,00:01:39
30     102.5,131.0
31     ```
32
33     - The `start,end` header is optional (auto-detected).
34     - Timestamps may be **decimal seconds** (`102.5`) or **colon-separated**
35       `MM:SS` / `HH:MM:SS` (`00:01:12`). Forms can be mixed.
36     - Blank lines and `#` comments are ignored.
37     - Malformed rows (bad timestamp, `start >= end`) fail loudly rather than being
38       silently skipped ? so labeling mistakes surface immediately.
39
40     Generate an empty template to fill in:
41
42     ```bash
43     python run_eval.py --scaffold --annotations rallies.csv
44     ```
45
46     ## 3. Running the evaluation
47
48     First make sure the video has been processed so its detections are in the DB
```

```

49 (the pipeline keys each video by its file MD5). Then:
50
51 ```bash
52 # Identify the video by file (its MD5 is computed for you)...
53 python run_eval.py --video "C:/videos/match.mp4" --annotations rallies.csv
54
55 # ...or by the stored video_id (MD5) directly:
56 python run_eval.py --video-id <md5> --annotations rallies.csv --stage both --iou 0.5
57 ```
58
59 Useful flags:
60
61 | Flag | Effect |
62 |-----|-----|
63 | `--stage {candidates,final,both}` | Which prediction stage(s) to score (default
64 `both`). |
65 | `--detector <name>` | Only score candidates from one detector (e.g. `yolo_hybrid`,
66 `tracknet_hybrid`). |
67 | `--iou <float>` | Match threshold (default `0.5`). |
68 | `--pr-curve` | Also report an IoU sweep (0.1?0.9). |
69 | `--show-failures` | List missed (FN) and spurious (FP) rallies with timestamps. |
70 | `--json <path>` | Write a machine-readable report. |
71 | `--db <path>` | Override the DB path (default: `output/<output.db_name>` from
72 `config.json`). |
73
74 ## 4. Reading the output
75
76 ```
77 =====
78 Rally evaluation ? video 'demo'
79 Ground-truth rallies: 3 | IoU threshold: 0.5
80 =====
81
82 Stage      Pred   TP   FP   FN   Prec   Rec   F1   mIoU  dStart  dEnd
83 -----
84 candidates    3    3    0    0   1.00   1.00   1.00   1.00   0.0s   0.0s
85 final          2    2    0    1   1.00   0.67   0.80   0.95   0.2s   0.5s
86 =====
87 [final] missed rallies (false negatives): 1
88   - 00:01:20?00:01:32
89 ```
90
91 - **TP / FP / FN** ? true positives (matched), false positives (spurious
92 predictions), false negatives (missed rallies).
93 - **Prec / Rec / F1** ? precision, recall, and their harmonic mean.
94 - **mIoU** ? mean temporal IoU over matched pairs (how tightly matched rallies
95 overlap their ground truth).
96 - **dStart / dEnd** ? mean absolute boundary error in seconds (are rallies
97 detected but trimmed in the wrong place?).
98
99 ### The key diagnostic: candidates vs. final
100
101 Scoring both stages tells you where to spend effort:
102
103 - **Candidates miss rallies the GT has** ? the **detector** is the bottleneck.
104 Improve the shuttle model (e.g. fine-tune TrackNet/WASB with frame-level
105 labels) or relax detection thresholds.
106 - **Candidates catch them but `final` is wrong/trimmed** ? the **segmentation /
107 AI handoff** is the bottleneck. Tune thresholds (`velocity_threshold`,
108 `dead_time_merge_gap`, `min_action_window_duration`) or the AI prompt ? no new
109 labels needed.
110
111 In the example above, detection is perfect (candidates: recall 1.0) but the AI
112 handoff dropped one rally (final: recall 0.67) ? so the fix is in segmentation,
113 not detection.

```

```

111     ## 5. How matching works
112
113     Predictions and ground truth are sets of `[start, end)` intervals. A prediction
114     **matches** a ground-truth rally when their **temporal IoU** (overlap ÷ union) is
115     at least the threshold. Matching is **greedy and one-to-one** ? highest-IoU pairs
116     are claimed first, and each prediction/GT is used at most once, so a single
117     detection can't be credited for two rallies. Unmatched predictions are false
118     positives; unmatched ground-truth rallies are false negatives.
119
120     ## 6. Code map
121
122     | File | Role |
123     |-----|-----|
124     | `backend/eval/metrics.py` | Pure interval math: IoU, matching, P/R/F1, boundary error.
125     | |
126     | `backend/eval/annotations.py` | Load/validate the ground-truth CSV. |
127     | `backend/eval/harness.py` | Read predictions from the DB, score, format the report. |
128     | `run_eval.py` | CLI entry point. |
129     | `tests/test_eval*.py` | Deterministic tests for all of the above. |
130
131     ## 7. Troubleshooting
132
133     | Symptom | Cause & fix |
134     |-----|-----|
135     | `ERROR: database not found at output/...` | The harness only reads stored predictions
136     ? it does **not** run the pipeline. **Process the video first** (e.g.
137     `scripts/run_process_and_stitch.py`) so detections land in the DB, then re-run. |
138     | Every rally shows as a miss (FN), `Pred` is 0 | This video hasn't been processed, or
139     you used the wrong `video_id`. The pipeline keys videos by **file MD5** ? pass the **same** file
140     via `--video`, or the exact stored `--video-id`. |
141     | `video_id` doesn't match although it's "the same" video | Re-encoding/trimming changes
142     the bytes, hence the MD5. Use the exact file that was processed, or look up the stored id (it's
143     printed during processing). |
144     | `candidates` stage is empty but `final` has data (or vice-versa) | Depends on which
145     segmenter ran. `gemini` writes only `final` (no candidates); `yolo_hybrid`/`tracknet_hybrid`
146     write both. Use `--stage final` for the gemini segmenter, and `--detector <name>` to scope
147     candidates. |
148     | Recall looks low only at high `--iou` | Rallies **are** found but boundaries are loose ?
149     check `dStart`/`dEnd`. Tune `stitching` padding or the segmentation thresholds rather than the
150     detector. |
151     | `ValueError: start (...) >= end (...)` while loading annotations | A bad row in your
152     CSV. Timestamps must be increasing per rally and use seconds or `HH:MM:SS`. |
153
154     Run `python run_eval.py -h` for a quick reminder of this prerequisite and all flags.

```

6. user (2026-07-06T20:14:06.777Z)

```

1     # Data in GCS ? the corpus source-of-truth map
2
3     > **Scope:** this map covers the **training corpus only** (`gs://khelsutra-rally-
4     corpus`). **User-uploaded
5     > videos are a SEPARATE, already-shipped surface** ? `POST /api/assets` + the per-
6     request `medium_uri` +
7     > `GET /api/capabilities.storage` (see `khelsutra/docs/STUDIO_MEDIA_LIFECYCLE.md`) ? and
8     must stay physically
9     > isolated from the corpus bucket (the **D1** compliance boundary: consent / no-minors /
10    not-training holds at
11    > the storage layer). The deploy layer's bucket names now live in one place
12    (`deploy/gcp/buckets.env`).
13
14    **Status:** LIVE / in-progress (migration not complete). **Owner-gated** decisions in
15    §7.
16
17    **Last inventoried:** 2026-06-22 against `gs://khelsutra-rally-corpus` (§5 gap table
18    reconciled 2026-07-02 to the post-Phase-1 state ? it had kept showing the pre-Phase-1 gaps,

```

contradicting \$7; Roora pilot non-video archive, Roora distilled golden-label fixtures, and six owned-source MD5s reconciled 2026-07-02).

11 **Goal (owner, 2026-06-22):** run **all** training + inference from the cloud with **all data in GCS**,

12 so the local Windows box (and the intermittently-mounted external drive) is **no longer a blocker**.

13

14 > **? Other sessions: this is where the data lives.** Pull from `gs://khelsutra-rally-corpus`, not

15 > from anyone's `C:` / `F:` / Google Drive. If something you need isn't in the bucket yet, it's in the

16 > **gap table (\$5)** ? flag it / migrate it, don't hard-code a local path.

17

18 This doc is the **data map**. The *machine* runbooks are
 SETUP_NEW_MACHINE.md

19 (spun up a box) and [deploy/gcp/README.md](../deploy/gcp/README.md) (provision GPU + train/infer).

20 The cloud engine itself was built by the (archived) DECOUPLED_COMPUTE project

21 (`docs/archives/past\_projects/DECOUPLED\_COMPUTE/`); this doc does **not** redesign it ? it records

22 *where the bytes are* and *what's still missing*.

23

24 ---

25

26 ## 1. The bucket

27

28 | | |

29 |---|---|

30 | Bucket | **`gs://khelsutra-rally-corpus`** (single corpus bucket) |

31 | Project | `gen-lang-client-0306026826` (billing *Khelsutra Billing*) |

32 | Region | `asia-south1` (Mumbai), UBLA, co-located with the GPU VMs |

33 | Auth | SA `rally-worker@gen-lang-client-0306026826.iam.gserviceaccount.com`
 (`roles/storage.objectAdmin`); key at `~/gcp/rally-worker-sa.json` (local only, **never committed**). Interactive: `gcloud auth login`. |

34 | Size today | ~2.06 GB (mostly past dry-run reels/outputs ? **not** the full corpus) |

35

36 The engine is **backend-agnostic by URI** (`backend/storage/ref.py`): `gs://` (or `gs://`),

37 `s3://`, `gdrive://`, or a bare local path. Pick the backend with **config + creds, no code change**.

38

39 ---

40

41 ## 2. Canonical layout (the target contract)

42

43 The cloud worker's input/output contract (`backend/worker/\_\_main\_\_.py`; lineage: `DECOUPLED\_COMPUTE/04`):

44

45

46 | Prefix | Holds | Who reads/writes it |

47 |---|---|---|

48 | `videos/` | source videos / proxies (<clip>.mp4) | **input** to `worker infer` & `worker train` |

49 | `labels/` | golden rally labels (<clip>.rallies.csv) | **input** ground truth to `worker train` + all eval |

50 | `trajectories/` | cached WASB trajectory CSVs | ? **eval/litmus input** (GPU-free), consumed via `manifests/golden.json`. **NOT** read by `worker train` (it synths or re-extracts) |

51 | `weights/<version>/` | trained Gen-0 bundles (`weights.json` + a run `manifest.json`) | **output** of `worker train` |

52 | `models/` | upstream pretrained weights (`wasb\_badminton\_best.pth.tar`) | **staged to a local path** (`\$WASB\_WEIGHTS` / `~/models/?`) for the detector ? not read from `gs://` in place |

53 | `outputs/<video\_id>/` | per-video inference outputs (intervals, etc.) | **output** of `worker infer` |

54 | `highlights/`, `promo/` | rendered reels + promo artifacts | output / demo |

55 | `archive/roora-pilot-2026-06/` | Roora pilot non-video deliverables, review notes, sidecar CSVs, trajectory exports, and a SHA-256 manifest | ****archive only****; raw videos remain owner-gated |

56

57 > ****Note on the manifest.**** `worker train` does ****not**** read a pre-staged corpus `manifest.json` ? it

58 > ***builds*** one on-box at train time from `labels/` (+`videos/`), extracting trajectories ****synthetically****

59 > by default****** (`--trajectory-source synth`, CPU-mock) or for real via opt-in `--trajectory-source`

60 > detector` (needs `videos/`). A `manifest.json` only appears as an ****output**** inside each

61 > `weights/<version>/` bundle. The ****eval/litmus**** path is what wants a stable, path-resolvable manifest

62 > (§7.3) ? and `manifests/golden.json` is currently consumed by ****no code yet**** (eval is local-FS only;

63 > see Known blockers).

64

65 ---

66

67 ## 3. What's actually in the bucket today (2026-06-22)

68

69 > ****Phase-1 landed 2026-06-22:**** added `trajectories/` (15), 15 canonical `labels/<date>\_<name>.rallies.csv`

70 > (= original-6 + `GX010137`/`GX010141` + 7 new; the 7 old `\_proxy`/bare dupes were deleted), and

71 > `manifests/golden.json`. The table below is the ****pre-Phase-1**** snapshot; the remaining gap is ****videos****

72 > (14 of 15) ? see §6/§7b/§8.

73

74	Prefix	Present	Drift / notes
75	--- --- ---		
76	`labels/`	7 files: `GX010128`, `mahadevpura-1`, + `Badminton BXH 2_proxy`, `Boxhill Doubles_proxy`, `GX020094_proxy`, `GX030094_proxy`, `mahadevpura-2_proxy`	? **inconsistent naming** (`_proxy` on some, not others); ? **missing** the original-6 labels + `GX010137`, `GX010141`
77	`videos/`	3: `GX030094_proxy.mp4`, `mahadevpura-1.MP4`, `mahadevpura_smoke_180s.mp4`	? corpus videos/proxies almost entirely absent
78	`trajectories/`	? ?	? **does not exist** (every trajectory is local-only)
79	`models/`	`wasb_badminton_best.pth.tar`	? (? WASB-C3 weight-provenance still an open legal item before any external share)
80	`weights/`	`0.1.0-l4/`	? first real L4 train bundle
81	`outputs/`	one `<<video_id>/`, `parity/`	? dry-run outputs
82	`highlights/`, `promo/`	reels + promo	?

83

84 ****Drift summary:**** the bucket has extra prefixes (`highlights/`, `promo/`, `models/` alongside

85 `weights/`) and ****inconsistent `labels/` naming****, and is ****missing the bulk of the corpus inputs****

86 (`videos/`, all `trajectories/`, most `labels/`). It reflects past ***dry runs***, not a complete corpus.

87

88 ****Roora pilot archive addendum (2026-07-02):**** the non-video pilot record is now backed up under

89 `gs://khelsutra-rally-corpus/archive/roora-pilot-2026-06/`. The prefix contains 740 objects total:

90 739 manifest entries (5,731,641 bytes) plus `manifest.json`. Verification on 2026-07-02 found

91 ****0 video objects**** under the prefix. The manifest schema is `khelsutra.roora\_pilot\_archive.v1` and

92 records relative local source paths, destination `gs://` URIs, byte sizes, and SHA-256 hashes. Raw

93 pilot videos under `Annotation Videos/` were deliberately excluded. The distilled human-corrected

94 Roora goldens are also git-versioned in `sports-data-collector/data/golden\_labels/` as

```

of
95     `sports-data-collector#56`.
96
97     ---
98
99     ## 4. How the cloud consumes it
100
101     Provisioning + the verified commands live in
[`deploy/gcp/README.md`](../deploy/gcp/README.md) §3?§5.
102     The data-facing entry points (`backend/worker/__main__.py`):
103
104     ```bash
105     # config (no secrets in config; auth via GOOGLE_APPLICATION_CREDENTIALS):
106     #   { "storage": { "backend": "local", "gcs": { "project": "gen-lang-client-0306026826"
} } }
107
108     # INFERENCE on one cloud video -> outputs/<video_id>/ in the bucket
109     python -m backend.worker infer \
110         --video-uri gcs://khelsutra-rally-corpus/videos/<clip>.mp4 \
111         --out-uri   gcs://khelsutra-rally-corpus --segmenter wasb_hybrid --config config.json
112
113     # TRAIN from the bucket (reads labels/ [+ videos/], extracts trajectories, builds
manifest on-box,
114     # runs training.gen0.harness, pushes weights/<version>/):
115     RALLY_DETECTOR_IMPL=native python -m backend.worker train \
116         --corpus-uri gcs://khelsutra-rally-corpus --out-uri gcs://khelsutra-rally-corpus \
117         --version <v> --trajectory-source detector --segmenter wasb_hybrid --config
config.json
118     ```
119
120     A bare/relative path resolves under a configured `input_root`/`input_backend`
121     (`backend/storage/ref.py::resolve_input_uri`), so workflows can be written bucket-
relative.
122
123     ---
124
125     ## 5. The gap ? what's still local-only (why the box is still a blocker)
126
127     > **Reconciled 2026-07-02.** Phase-1 corpus *metadata* (trajectories + labels +
manifest) **landed in
128     > the bucket 2026-06-22** ? see §6 (Phase 1) and §7 decisions #2/#3 (LOCKED/done). This
table previously
129     > still showed those as ? and contradicted §7/§7b. Post-Phase-1 the remaining gaps are
the **videos**
130     > (Phase 2) and the **Phase-3 eval-on-GCS plumbing** (§6 Phase 3 ? the eval drivers
still read local
131     > paths with plain `open()`, so bucket-resident metadata isn't consumed by eval yet).
132
133     | Data | Where it is now | Size | In bucket? | Needed for |
134     |---|---|---|---|---|
135     | WASB trajectory CSVs | bucket *** local working copy (`?\Annotation
Setup\?\Trajectories\`) | **17.5 MB** | ? canonical `trajectories/<name>.csv` (2026-06-22, §7.2)
| eval/litmus (cached). *(No `cached` train source exists ? train synths or detector-extracts.)*
|
136     | golden manifest | `gs://?/manifests/golden.json` (§7.3); the **local** litmus still
reads `output/human_lovo_manifest.json` | 4.5 KB | ? (2026-06-22, §7.3) ? flipping the litmus to
read it is Phase 3 | eval corpus definition |
137     | original-6 labels | bucket **** F: archive | ~12 KB | ? among the 15 canonical
`labels/<date>_<name>.rallies.csv` (2026-06-22, §7.1) | train + eval |
138     | 7 new-clip labels | bucket **** repo `output/*.rallies.csv` (gitignored) **** F:
archive | ~12 KB | ? ditto (naming canonicalized 2026-06-22; the stale `_proxy`/bare dupes
MD5-verified + deleted, §7) | train + eval |
139     | source videos / proxies | F: archive + Google Drive golden folder; six owned-source
MD5s in `sports-data-collector/data/golden_source_videos.json` | **100s of GB** | ? mostly not
migrated ? **the remaining large gap** (Phase 2; see §6) | infer + train (real WASB extraction)

```

```

|
140 | Roora pilot non-video record | `Annotation Setup\Archive\Roora Pilot Review - 6 June
2026` + root sidecar CSVs + `Trajectories\roora`; distilled goldens also in `sports-data-
collector/data/golden_labels/` | **5.7 MB** | ? `archive/roora-pilot-2026-06/` (2026-07-02
manifest + 0 videos); small goldens git-versioned in `sports-data-collector#56` | vendor-pilot
archive / loss prevention |
141 | Roora pilot raw videos | `Annotation Setup\Archive\Roora Pilot Review - 6 June
2026\Annotation Videos` | **4.8 GB** | ? owner-gated; not uploaded by the metadata archive pass
| vendor-pilot provenance |
142
143 > Post-Phase-1 the corpus *metadata* is **in the bucket**; what still blocks GPU-free
cloud eval is the
144 > **Phase-3 plumbing** ($6 ? the eval drivers can't fetch `gs://` yet). The **videos**
are the only large
145 > migration left (Phase 2), and they already exist on **Google Drive** (the golden-
corpus folder), so they
146 > can go **Drive ? bucket from a cloud box**, keeping the local machine out of the path
entirely.
147
148 ---
149
150 ## 6. Migration plan (close the gap)
151
152 **Phase 1 ? corpus metadata (?25 MB, do first, removes the eval blocker). ? DONE
2026-06-22**
153 (see $7 decisions #1/#2/#3 + $7 cleanup note; the sketch below is kept as the historical
plan). Stage all
154 trajectories, all labels, and a **bucket-relative eval manifest** into `gs://khelsutra-
rally-corpus`.
155 Sketch:
156
157 ```bash
158 # trajectories (decision-free; absent today)
159 gcloud storage cp "C:\?\Annotation Setup\Trajectories\*.csv" gs://khelsutra-rally-
corpus/trajectories/
160 # labels (after the $7.1 naming decision) -> labels/<canonical-name>.rallies.csv
161 # eval manifest with gs:// (or bucket-relative) paths -> a stable manifest object ($7.3)
162 ```
163
164 **Phase 2 ? videos ? `videos/<name>.mp4`** (bucket name = the eval `name`, consistent
with
165 `trajectories/` + `labels/`). Sources (resolved 2026-06-22 ? see
`scratch/copy_videos_to_gcs.ps1`):
166
167 | Source | Clips | Size |
168 |---|---|---|
169 | Drive golden folder `~/Golden Label Videos Request - June 15/[Done] Video N/`
(**proxies**) | `mahadevpura_1/2`, `GX010128`, `GX020094`, `GX030094`, `Badminton_BXH_2`,
`Boxhill_Doubles` (7) | ~13.5 GB |
170 | F: archive (**full MP4s** ? proxy-first before upload; owner-box root verified as
`F:\[Khelsutra] GoPro Hero Black 12\GoPro Recorded Videos`) | `GX010137`, `GX010141`, +
original-6 (`adarsh_avi_singles`, `kushagra_singles`, `largetest_doubles`, `gbaaddy`,
`testlarge_short`, `mahadevpura_singles`) (8) | ~50 GB |
171
172 **Two ways to run it:**
173 - **Decoupled (preferred ? no local box):** `rclone` on a cloud VM with a Google-Drive
remote ?
174 `rclone copy gdrive:"<?/[Done] Video N>/<proxy>.mp4" :gcs:khelsutra-rally-
corpus/videos/<name>.mp4`
175 (one-time owner step: `rclone config` ? `gdrive` remote OAuth). Or `gdown <file-id>`
on the box for
176 the shared folder. This is the path that actually removes the local box.
177 - **Interim (today, owner box ? streams Drive/F: ? local ? GCS):** `pwsch
scratch/copy_videos_to_gcs.ps1
178 -Execute` (dry-run without `-Execute`; `-Only <name>` for one; skip-existing). Uses

```



```

local bandwidth.
179
180     **Priority:** the 7 Drive **proxies** first (already 1080p, sufficient for detection).
The 8 large F:
181     full MP4s should be **proxied to 1080p first** (cheaper storage + faster cloud decode)
before upload.
182     ? Two pre-existing non-canonical video objects (`videos/GX030094_proxy.mp4`,
`videos/mahadevpura-1.MP4`)
183     are superseded once `videos/GX030094.mp4` / `videos/mahadevpura_1.mp4` land ? verify-
dupe-then-delete
184     (same as the label cleanup). Track per-clip in $8.
185
186     **Phase 3 ? make eval/litmus read GCS (real code, NOT a config flip).** The eval drivers
187     (`served_gate_a.evaluate_video`, `calibrate_wasb.load_golden_gts`,
`TrackNetRunner.parse_trajectory_csv`)
188     open `spec["trajectory"]`/`spec["labels"]` with plain `open()` and gate on
`os.path.isfile()` ? they
189     have **no storage-layer fetch**, so a `gs://` manifest cannot be consumed today. Phase-3
must thread
190     `storage.fetch`/`resolve_input_uri` through those paths first. Then: re-probe
`fps`/`frame_width` from
191     the bucket videos (don't trust the copied values ? a wrong fps silently mis-aligns every
label), and
192     **re-measure the served Gate-A litmus on the GCS 15-clip set as a [[launch-report-
convention]] event**
193     (dual-set RAW+golden no-regression, not a silent eval-set swap ? see PR #236).
194
195     ---
196
197     ## 7. Decisions
198
199     1. **labels/` naming ? LOCKED (owner, 2026-06-22):** `labels/<YYYY-MM-
DD>_<name>.rallies.csv`, where
200     `<name>` is the eval/manifest clip name and `<YYYY-MM-DD>` is the label's **authoring
date** =
201     the *earliest* mtime across all known copies (C: working copy + F: archive ? robust
to copy-resets;
202     a restored C: copy shows today, the F: archive keeps the real date). Fixes the
`_proxy` inconsistency
203     and disambiguates recycled GoPro names by date. *(Future, when videos are in the
bucket: optionally
204     add a `video_id`/MD5 sidecar for the strongest keying ? deferred, needs hashing each
video.)*
205     2. **Cache `trajectories/` ? YES, done 2026-06-22.** `trajectories/<name>.csv` cached so
cloud/CI eval +
206     the served Gate-A litmus run **GPU-free** (training can still regenerate via
`--trajectory-source detector`).
207     3. **Eval manifest in GCS ? done 2026-06-22.** `gs://khealsutra-rally-
corpus/manifests/golden.json`
208     (15 clips, `gs://` trajectory + label paths). The **local** litmus keeps its local-
path manifest;
209     flipping the litmus to read the GCS manifest is Phase 3.
210     4. **Originals migration scope ? OPEN** ? proxies-only vs full 4K, and how much of the
15-clip golden
211     set vs the wider archive. Source = Google Drive ? bucket from a cloud box (not the
local box).
212     5. **Roora pilot non-video archive ? done 2026-07-02.** `archive/roora-pilot-2026-06/`
backs up
213     CVAT/vendor deliverables, feedback/review notes, root sidecar CSVs, and
`Trajectories/roora`
214     sidecars with SHA-256 hashes in `manifest.json`; the small human-corrected
badminton/pickleball/
215     table-tennis goldens are git-versioned in `sports-data-collector#56`; raw videos stay
excluded
216     pending owner upload scope.

```

```

217     6. **Owned-source MD5 record ? done 2026-07-02.** The six original owned golden MP4s on
`F:` were
218     MD5-hashed and the collector registry paths were refreshed in `sports-data-
collector#60`; the
219     durable manifest is `sports-data-collector/data/golden_source_videos.json`. Uploading
those large
220     originals (or generated proxies) remains the Phase-2 owner-gated decision.
221
222     > Phase-1 was run via a dry-run planner (resolve every clip's trajectory + label, stamp
the label's
223     > authoring date, build the `gs://` manifest) ? `gcloud storage cp` (planner kept in
gitignored
224     > `scratch/`; promote it to `scripts/` if corpus-sync recurs).
225     > ? **Cleanup DONE 2026-06-22:** the 7 pre-existing inconsistent label objects
(`*_proxy.rallies.csv`,
226     > bare `GX010128`/`mahadevpura-1`) were confirmed **byte-identical (MD5)** to the new
canonical names
227     > (content also on C:/F:) and **deleted** ? `labels/` now holds exactly the 15 canonical
objects.
228
229     > ? **gcloud gotcha (Phase-2):** `gcloud storage cp` treats `[...]` as a glob, so source
paths under
230     > `[Done] Video N/` (Drive) or `[Khelsutra] ?` (F:) fail with "matched no objects." The
owner-box
231     > script (`scratch/copy_videos_to_gcs.ps1`) stages such paths through a bracket-free
temp file; the
232     > `rclone` cloud path has no such issue.
233
234     ---
235
236     ## 7b. Known blockers ? must-fix before the DoD (2026-06-22 adversarial review)
237
238     1. **? RESOLVED 2026-06-22 ([#319](https://github.com/Khelsutra/badminton-highlight-
indexer/pull/319))
239     ? dated labels broke `worker train --trajectory-source detector`. **History:**
`cmd_train` derived the
240     join key as `fname.replace(".rallies.csv","")` ? `2026-06-22_GX020094`, then looked
up
241     `video_by_stem.get(name)` against bare video stems (`GX020094`), so every clip was
skipped and train
242     aborted with `<2 videos`. **Fix:** the label?clip-name derivation is now the helper
`_label_clip_name()`
243     in [backend/worker/`__main__.py`](../backend/worker/`__main__.py`), which strips both the
244     `.rallies.csv`/`.csv` suffix AND a leading `^d{4}-d{2}-d{2}_` prefix, so a dated
label and a legacy
245     bare label both resolve to the bare `videos/<name>` stem (preserves the locked human-
readable naming).
246     Tested: parametrized helper cases + a dated-label?bare-video `cmd_train` join
regression + per-label
247     skip-branch coverage. *(`synth` train + all eval were unaffected ? they don't join
videos.)*
248     2. **Phase-3 eval-on-GCS is net-new plumbing, not a flag** ? see §6 Phase-3. Eval is
local-FS only.
249     3. **`manifests/golden.json` is consumed by zero code today** (train ignores it; eval
can't fetch
250     `gs://`) ? it's correct-but-dormant until blocker #2 lands. Add a tiny manifest-
validator + a CI job
251     that authenticates to GCS and runs the litmus against it, or the DoD ("CI reads only
from `gs://`")
252     stays unverifiable.
253     4. **Gating infra (owner-only, not done):** GPU `GPUS_ALL_REGIONS` quota is still **0**
(no VM ever
254     created ? `deploy/gcp/README` §2); cloud-box secrets (SA JSON + rotated
`GEMINI_API_KEY` + WASB
255     weights staged from `models/`) aren't set up persistently. No real cloud train/infer

```

```

runs until these.
256     5. **Reproducibility:** the migration + copy helpers live in gitignored `scratch/`
257     (`migrate_golden_to_gcs.py`, `copy_videos_to_gcs.ps1`) ? promote to `scripts/` so
Phase-2/re-syncs
258     are repeatable from the repo, and add a referenced **proxy-generation** step for the
8 large F: originals.
259
260     ## 8. Progress tracker
261
262     | Item | Status |
263     |---|---|
264     | Bucket + SA + provisioning (deploy/gcp) | ? built |
265     | Cloud worker `infer`/`train` by URI | ? built (`backend/worker`) |
266     | This data-map doc | ? 2026-06-22 |
267     | Phase 1: trajectories ? `trajectories/` (15) | ? 2026-06-22 |
268     | Phase 1: labels reconciled ? `labels/<date>_<name>` (15) | ? 2026-06-22 |
269     | Phase 1: eval manifest ? `manifests/golden.json` | ? 2026-06-22 |
270     | Phase 1: remove 7 stale `_proxy`/bare label objects | ? 2026-06-22 (MD5-verified
dupes, deleted) |
271     | Blocker $7b#1: dated labels join in `worker train --trajectory-source detector` | ?
2026-06-22 ([#319](https://github.com/Khelsutra/badminton-highlight-indexer/pull/319)) ?
`_label_clip_name` strips the ISO-date prefix |
272     | Roora pilot non-video archive ? `archive/roora-pilot-2026-06/` | ? 2026-07-02 ? 739
manifest entries + `manifest.json`; verified 740 objects and 0 video objects in the prefix |
273     | Roora pilot distilled golden labels | ? 2026-07-02 (`sports-data-collector#56`) ?
human-corrected badminton + pickleball/table-tennis goldens are git-versioned with tests |
274     | Six owned golden source MP4 hashes | ? 2026-07-02 (`sports-data-collector#60`) ?
MD5/path manifest committed; upload still owner-gated |
275     | Phase 2: corpus videos/proxies ? `videos/<name>.mp4` | ? mapping + script done;
**1/15** demo'd (`mahadevpura_1`); 14 remain ? run `rclone` on a cloud box (preferred) or the
owner-box script |
276     | Phase 2: proxy the 8 large F: full-MP4 clips before upload | ? |
277     | Phase 3: eval/litmus default to the GCS manifest | ? |
278
279     **Definition of done:** a fresh GPU box (or CI) can run inference, the eval suite, and
Gen-0 training
280     reading **only** from `gs://khelsutra-rally-corpus` ? no `C:` / `F:` / local-Drive path
required.
281
282     ---
283
284     ## Related
285     - [`SETUP_NEW_MACHINE.md`](SETUP_NEW_MACHINE.md) ? spin up a box (storage backends in
$7).
286     - [`deploy/gcp/README.md`](../deploy/gcp/README.md) ? GCP GPU provisioning + train/infer
($5).
287     - [`VIDEO_LOCALITY_MODEL.md`](VIDEO_LOCALITY_MODEL.md) .
[`PER_VIDEO_WORKSPACE.md`](PER_VIDEO_WORKSPACE.md) ? locality tiers + `outputs/<video_id>/`.
288     - `docs/archives/past_projects/DECOUPLED_COMPUTE/` ? the cloud-engine design lineage.
289

```

7. user (2026-07-06T20:14:07.059Z)

```

1     # R2 ? Task-faithful rally-boundary eval metric (tolerance-match + over-seg guard)
2
3     **Status:** **IMPLEMENTED (augmenting), 2026-06-17** ? owner-reviewed + decisions locked
2026-06-17;
4     `backend/eval/tolerance_metrics.py` + `rally_seg_eval` reporting landed (default-
additive; tIoU stays the
5     gate pending the ablation experiment). Gates the rally-end work (R4) and the dual-eval-
set regression gate.
6     A focused increment of the rally-quality project
([RALLY_QUALITY_RESEARCH.md](RALLY_QUALITY_RESEARCH.md) $7;
7     the day-1 validation that motivates it:
[REAL_FOOTAGE_VALIDATION_REPORT.md](REAL_FOOTAGE_VALIDATION_REPORT.md);

```

```

8     the directive: `eval-boundary-tolerance` + `eval-dual-set-regression`).
9
10    > **Implementation refinement (2026-06-17):** §2.3.2's boundary-error distribution is
computed over
11    > **unconditional best-overlap** matches, NOT the within-? 1:1 matches ? a ?-gated set
is survivorship-biased
12    > (loose ends fail the gate and vanish, hiding the very end-deficit the diagnostic
exists to surface; that
13    > within-? deficit instead shows up as low `tol_recall`). Matches the day-1 validation
method. Measured on the
14    > n=9 corpus: **|?start| 1.25 s vs |?end| 2.19 s** ? "start?solved, end=the gap",
confirmed.
15
16    ---
17
18    ## 0. TL;DR
19    `tIoU-F1@0.5` is the wrong ruler for rally detection, in **both** directions: it over-
penalizes the
20    **intrinsically fuzzy rally START** (the ~1?1.5 s service window), and its overlap
matching is murky about
21    **over-production**. R2 replaces the headline with a **tolerance-match metric under
strict 1:1 (Hungarian)
22    assignment**, reported as a **decomposition** ? `P@? / R@? / F1@?` + **boundary-error
distributions split
23    into start/end** + an **over-segmentation guard** + a **?-sweep** ? not a single number.
24
25    **Honest, slightly uncomfortable consequence (measured, see §3):** under R2 the local
detector's headline
26    number *drops* and the gap to Gemini *widens* vs what tIoU showed ? because R2
faithfully charges for the
27    1.5× over-production and the loose ends that tIoU's overlap matching partly absorbed.
That is the point:
28    R2 splits "behind Gemini" into its two real causes ? **a) over-production [precision]**
and **b) loose
29    rally-ends [the R4 lever]** ? and forgives only the inherent start fuzziness. It is a
more honest ruler,
30    not a flattering one.
31
32    ---
33
34    ## 1. Why `tIoU@0.5` is the wrong ruler (grounded in the n=2 ground-truth clips)
35    1. **Over-penalizes the START.** For a 5.7 s rally, `tIoU ? 0.5` tolerates only ~±1.9 s
of boundary shift
36    (`? = D·(1??)/(1+?)`). The service window (stance ? racquet contact) alone is ~1?1.5
s ? even **Gemini's**
37    median |?start| is 0.98?1.46 s. So tIoU spends most of its budget on a boundary that
is **intrinsically**
38    undefined to ~1 s. (Owner directive: don't strictly align starts; tolerate ~±2 s.)
39    2. **Count parity is a deceptive proxy.** mahadevpura-2 hit 1.03× of Gemini's count yet
F1@0.5 = 0.150 ?
40    right count, wrong boundaries. A count- or overlap-based view hides this.
41    3. **It flips config rankings on noise.** The hard-cap was golden-neutral at n=6 (+0.006
n.s.) but "best" at
42    n=7 once a continuously-active clip entered ? a metric artifact at the 0.5 cliff, not
a quality change.
43    4. **The deficits we must see are start?end-symmetric.** Local's |?start| ? Gemini's;
the **entire** gap is
44    the rally-**END** (|?end| 2.1?3.1 s vs Gemini ~0.8?1.2 s). A single blended IoU
cannot show that ? and
45    directing R4 needs exactly that decomposition.
46
47    ---
48
49    ## 2. The metric
50

```

```

51     ### 2.1 Match rule (asymmetric tolerance ? forgive the start, expose the end)
52     A predicted window matches a golden rally iff **|?start| ? ?_start` AND `|?end| ?
?_end`**.
53     - **?_start` generous (~2.0 s)** ? absorbs the service-window/annotator start fuzziness
(the part that is
54         *not* a model deficit).
55     - **?_end` tighter (~1.5 s, but IAA-calibrated ? see §2.4)** ? the rally-end IS a sharp
physical event
56         (shuttle down/net/out), so a loose end *is* a real deficit we want the metric to
expose, not forgive.
57
58     > A *symmetric* ±2 s tolerance is tempting (the owner's stated bar) but the preview (§3)
shows it is **not**
59     > automatically more lenient than tIoU 0.5 and it blurs exactly the end-deficit we need
to see. Asymmetry is
60     > the point: relax the dimension that's intrinsically fuzzy (start), keep honest the one
that's a real event (end).
61
62     ### 2.2 Assignment ? strict 1:1 (Hungarian), never overlap/greedy
63     Match references to predictions by **global Hungarian assignment** (max total matches
under the tolerance
64     gate), one-to-one. This is what makes **over-production cost precision directly**: `P@?`
= matched / n_pred`
65     crashes when the detector emits 75 windows for 52 rallies. (tIoU's overlap matching let
extras "match
66     something"; the W1 over-segmentation pendulum was barely costed.)
67
68     ### 2.3 Report the DECOMPOSITION, never one number
69     1. **`P@?`, `R@?`, `F1@?`** ? precision (over-production), recall (did we find the rally
within tolerance), F1.
70     2. **Boundary-error distributions** ? median *** IQR** of signed *and* absolute
`?start`, `?end` over matched
71     pairs, **split**. This is the metric that says **"start solved, end open"* and aims
R4.
72     3. **Over-seg guard (mandatory):** `split_count` (preds covering one ref ? W1 over-seg)
and `merge_count`
73     (refs under one pred ? the baseline mega-window). The *original* rule was strict per-
video: **a promotion
74     may not increase either on ANY video.** **REFINED to a NET guard, 2026-06-17 (R1,
owner-approved)** ? the
75     strict rule is correct for an *incremental* cue but **mis-fires on a structural mega-
window?bounded change**
76     (splitting a mega-window necessarily lifts `split_count` from ~0, and the raw merge
*count* is non-monotonic:
77     one window hiding 40 rallies is `count=1, rate=1.0`; bounding it into pieces that
each glue 2 neighbours is
78     `count=9, rate=0.5` ? FEWER rallies hidden). The net guard scores over-seg on the
**normalized rates** as one
79     weighted cost per video (`w_merge·merge_rate + w_split·split_rate`, merges weighted
higher ? a merge hides a
80     rally = recall loss, a split only fragments one), and **passes iff (a) the corpus-
mean net cost does not rise
81     AND (b) no video's `merge_rate` regresses.** Implemented as
`tolerance_metrics.over_seg_guard` (it always
82     reports the strict verdict too; revert = read `strict_passes`). See
QUALITY_ITERATIONS.md "R1" for the measured
83     verdict (milestone: mean `merge_rate` 0.694±0.150, net PASS, strict BLOCK).
84     4. **?-sweep** (` ? {1.5, 2.0, 2.5, 3.0}`, symmetric for the trend; asymmetric pair for
the headline) ?
85     because the convention/IAA noise is large, a single ? misleads (§3). Report the
curve.
86     5. **`tIoU`/`mIoU` demoted to a SECONDARY trend-line diagnostic** (keeps the paper-
ablation backbone
87     comparable) ? **not the gate.**
88

```

```

89     ### 2.4 Choosing ? ? it MUST be calibrated to inter-annotator agreement
90     The preview (§3) shows even Gemini scores `tolF1@1.5 = 0.20` on mp2 ? i.e. on ~half
the rallies the
91     golden END differs from Gemini's by >1.5 s. That is plausibly label/convention noise,
not model error**:
92     we'd be measuring how two humans (or human-vs-Gemini) disagree on "when did the rally
end." Therefore
93     `?_end` must be ? the inter-annotator end-agreement**, which we do not yet know.
Action: a small IAA
94     study ? a 2nd labeler re-labels 2?3 already-golden videos; set `?_start` `?_end` to
~the 90th-percentile
95     pairwise human |?start|/?end|. Until then freeze provisional ?_start=2.0, ?_end=1.5
and always report the
96     sweep**, with the `? = D.(1??)/(1+?)` derivation recorded so the choice is auditable.
97
98     ---
99
100    ## 3. Preview on the two ground-truth clips (greedy-1:1 approximation; Hungarian in the
impl)
101
102    Symmetric-? sweep, F1@? (local-best = the per-clip best config: mp2 W1+W2+hardcap,
GX010128 W1+W2):
103
104    | clip / method | tolF1@1.5 | @2.0 | @2.5 | @3.0 | (tIoU-F1@0.5) |
105    |---|---|---|---|---|---|
106    | mahadevpura-2 / Gemini | 0.203 | 0.354 | 0.506 | 0.557 | 0.557 |
107    | mahadevpura-2 / local | 0.111 | 0.148 | 0.222 | 0.278 | 0.296 |
108    | GX010128 / Gemini | 0.615 | 0.769 | 0.813 | 0.857 | 0.813 |
109    | GX010128 / local | 0.189 | 0.315 | 0.346 | 0.362 | 0.457 |
110
111    Reading the preview (all honest):
112    - tolF1@~3.0 ? tIoU-F1@0.5 ? the sweep is a smooth, interpretable knob that
contains tIoU as a point;
113    good (continuity with the frozen baseline).
114    - The local?Gemini gap WIDENS under the tighter, 1:1 tolerance (GX010128 ratio 0.56
at tIoU0.5 ?
115    0.41 at ?2.0). R2 charges local for the 1.5× over-production (precision) and the
loose ends ? exactly
116    the deficits tIoU under-counted. This is the metric working, not local regressing.
117    - Even Gemini drops at small ? (mp2 0.557?0.203 from ?3.0?1.5) ? strong evidence the
golden END
118    convention carries >1.5 s of noise ? §2.4's IAA calibration is a prerequisite, not a
nicety.
119    - Boundary-error medians (the robust diagnostic, from the validation report): local
|?start| 0.78?1.74 s
120    (? Gemini 0.98?1.46) vs local |?end| 2.12?3.12 s (? Gemini 0.76?1.19) ? the headline
R4 signal.
121
122    ---
123
124    ## 4. Promotion / "trustworthiness" rule (small-n)
125    A candidate is promotable only if all hold:
126    1. Paired per-video sign-test sweep on F1@? is near-unanimous: 6/0 @ n=6
(p=0.031) or 7/0 @ n=7
127    (p=0.016) (8/0 @ n=8, p=0.008); 5/1 or 6/1 is NOT significant. Run on the
growing-golden set.
128    2. No over-seg regression ? the net over-seg guard passes (§2.3.3): the corpus-
mean weighted
129    merge/split-rate cost does not rise AND no video's `merge_rate` regresses. (Was the
strict per-video
130    count rule; refined to net on 2026-06-17 ? see §2.3.3.)
131    3. Effect size + bootstrap CI reported as an honesty band (expect it to span 0 at
this n; use it to
132    separate "real" from "directionally-clean-but-negligible," not as the pass/fail).
133    4. Stratify, don't blend ? report the continuously-active blob clip (mahadevpura-1)

```

```

as its own stratum
134     (its near-degenerate scores can swing a small-n mean).
135 5. **No regression on EITHER eval set** (raw + growing-golden) ? the dual-set
discipline.
136
137 ---
138
139 ## 5. Implementation plan (pure, offline, additive ? zero production risk)
140 - **`backend/eval/tolerance_metrics.py`** (pure, GPU-free, unit-testable):
141 - `match_1to1(preds, gts, ?_start, ?_end) -> (matches, P, R, F1)` ? Hungarian via
142 `scipy.optimize.linear_sum_assignment` if available, else a pure  $O(n^3)$  fallback
(corpus n is tiny).
143 - `boundary_error_report(preds, gts) -> {?start, ?end: signed/abs median+IQR}`
(computed over
144 unconditional best-overlap matches internally ? see the refinement note above).
145 - `over_seg_report(preds, gts) -> {split_count, merge_count}` (reuse the bipartite
logic from
146 `segmentation_metrics.merge_split_report`).
147 - **Wire into `backend/eval/rally_seg_eval.py`** ? emit the R2 block **alongside** tIoU
(don't remove tIoU);
148 add `--tau-start/--tau-end` (both landed; the ?-sweep helper
`tolerance_metrics.tolerance_sweep` exists
149 but is not yet surfaced as a `rally_seg_eval` CLI flag). Surface in the experiment
leaderboard.
150 - **`run_regression.py`** ? add the dual-set R2 aggregate + the over-seg-guard gate.
*(Status: the guard
151 primitive landed as `tolerance_metrics.over_seg_guard` and is wired into
`rally_seg_eval.compare()`/`--vs`
152 ? the natural home, since `rally_seg_eval` loads the golden manifest + trajectories.
The dual-set
153 `run_regression.py` aggregate is still a separate future item.)*
154 - **Tests** ? synthetic intervals: exact-match, start-fuzzy-within-?, end-loose-
beyond-?, over-production
155 (precision crash), the Hungarian-vs-greedy tie case, empty inputs.
156 - **Default-additive ? ablation-gated replacement:** R2 is reported next to tIoU
(augment). The *gate*
157 switches to R2 only after a dedicated **ablation experiment** on the golden corpus
shows R2 ranks candidates
158 at least as sensibly as tIoU ? so the promotion bar never changes silently mid-flight.
When tIoU is retired
159 as the gate it is **deprecated with documentation, not deleted** (kept as the
secondary trend-line + the
160 `QUALITY_ITERATIONS.md` paper-ablation backbone, with a dated note recording the
ablation evidence for the
161 switch). Owner decision, 2026-06-17 (§7).
162
163 ---
164
165 ## 6. Definition of done
166 - `tolerance_metrics.py` + tests merged (default-additive); `rally_seg_eval` reports the
R2 decomposition
167 next to tIoU (the ?-sweep helper exists in `tolerance_metrics` but is not yet wired
into the
168 `rally_seg_eval` report); the corpus (n?8) re-scored under R2 and recorded in
QUALITY_ITERATIONS.
169 - The IAA study (§2.4) scoped (it can land after R2's code; R2 ships with provisional ?
+ the sweep).
170 - `run_regression.py` carries the dual-set R2 gate + the per-video over-seg guard ? but
R2 **augments** tIoU
171 until the ablation experiment (§7) validates it as the gate; only then does the gate
switch.
172 - tIoU/mIoU retained as a secondary trend-line; when retired as the gate, **deprecated-
with-docs (kept for
173 posterity)** per the §7 owner decision ? a dated note records why and links the
ablation evidence.

```

```

174
175 ---
176
177 ## 7. Decisions
178 **Locked (recommended, design-level):** Hungarian 1:1 (not greedy/overlap); report the
decomposition + sweep,
179 never one number; tIoU demoted to secondary; R2 default-additive until validated;
boundary-error split
180 start/end is the R4-aiming diagnostic.
181
182 **Locked (owner review, 2026-06-17):**
183 - **?_end asymmetric (1.5 s) vs ?_start (2.0 s) ? ADOPTED.** Asymmetry is the point:
forgive the intrinsically
184 fuzzy start, keep the rally-end honest. Your "±2 s, not more" bar maps to
`?_start=2.0`; `?_end` stays the
185 tighter, IAA-calibrated knob (§2.4).
186 - **IAA study after R2 code ? ADOPTED.** R2 ships now with provisional `?_start=2.0 /
?_end=1.5` + the full
187 sweep; the IAA study lands next to fix ?. R2's code is not blocked on it.
188 - **Gate: AUGMENT now, then ablation-gate the eventual replacement, then deprecate-with-
docs ? ADOPTED.**
189 R2 is reported **alongside** tIoU (augment) until validated. tIoU is **not** swapped
out silently: before
190 R2 becomes the promotion gate, run a dedicated **ablation experiment** on the golden
corpus showing R2
191 ranks candidates at least as sensibly as tIoU (and exposes the deficits tIoU hid).
When tIoU is finally
192 retired as the gate, it is **formally deprecated with documentation ? kept for
posterity, never deleted**:
193 it remains a secondary trend-line and the paper-ablation backbone
('QUALITY_ITERATIONS.md'), with a dated
194 note recording *why* it was superseded and the ablation evidence behind the switch.
(Consistent with the
195 repo's "retain weak/superseded signals; numbers inform, foundations decide" norm.)
196

```

8. assistant (2026-07-06T20:14:09.574Z)

Now let me read the three files in `docs\data\_pipeline\`. Let me first verify their exact paths.

9. user (2026-07-06T20:14:10.087Z)

```

docs\data_pipeline\GOLDEN_DATA_SHARING.md
docs\data_pipeline\GOLDEN_SET_IMPLEMENTATION.md
docs\data_pipeline\GOLDEN_SET_PHASE1_VIDEOS.md
docs\data_pipeline\README.md
docs\data_pipeline\VIDEO_RECORDING_GUIDELINES.md
docs\data_pipeline\GOLDEN_VIDEOS.md

```

10. assistant (2026-07-06T20:14:11.973Z)

All three requested files exist. Let me read them fully.

11. user (2026-07-06T20:14:13.479Z)

```

1 # Golden-labeled video registry
2
3 **Purpose.** The single source of truth for every **golden** (human rally-boundary-
labeled) video ? what
4 it is, where the source proxy lives, how many rallies, and whether it's ingested into
the eval corpus.
5 Use this to **isolate the golden videos into a dedicated location** and to keep the
growing golden set
6 auditable. Updated whenever a new video is labeled + ingested.

```



```

7
8     > **This is the GROWING GOLDEN-finalized set** in the dual-eval-set discipline (see
9     > [QUALITY_ITERATIONS.md](../QUALITY_ITERATIONS.md), the 2026-06-16 capstone + `eval-
dual-set-regression`):
10    > a **RAW eval set** (currently-bad videos) + this **golden-finalized eval set**
(promoted when at-par) +
11    > a separate **golden TRAINING set**. Every quality launch must not regress on either
eval set.
12    > Privacy-safe, video-free *shipping* of these =
[GOLDEN_REGRESSION_FIXTURES.md](../GOLDEN_REGRESSION_FIXTURES.md).
13
14    ## Contracts / conventions
15    - **Labels:** rally-annotator CSV
(`rally_number,start_time,end_time,ending_reason,sport,shots_count`,
16    decimal seconds) ? ingested by `backend/eval/gt_loader.py` (reads
`start_time`/`end_time`).
17    - **Alignment is mandatory before trusting any score.** The label timebase MUST match
the windowed proxy
18    (same fps + frame count). Verify with `ffprobe` (the `largetest_doubles` near-miss is
why) ? ideally
19    annotate the *exact* proxy the detector ran on (then `video_id` = MD5 matches and
alignment is free).
20    - **`video_id`** = MD5 of the proxy (`backend/utils/hashing.py::compute_video_id`); it
keys the WASB
21    trajectory cache, the Gemini reference rows, and the golden features.
22    - **Ingestion** = copy the trajectory + labels to durable paths, write
`output/<name>_golden_features.json`
23    (`gts` + candidates), append a row to `output/human_lovo_manifest.json`.
(`scratch/at_par.py` scores a
24    single video vs Gemini + golden; `backend/eval/rally_seg_eval.py` runs the whole
corpus.)
25    - **Path smoke check:** before running corpus evals, run
26    `python -m backend.eval.golden_manifest --manifest output/human_lovo_manifest.json
--min-count 15`.
27    The checker rebases stale absolute paths against sibling `Annotation Setup` folders
and `output/`, and
28    exits non-zero if any required trajectory or labels file is missing. Add `--write
output/human_lovo_manifest.local.json`
29    when a resolved local manifest artifact is useful.
30    - **Guardrails:** owned footage only; **minors excluded**; labels never derived from a
paid LLM (Gemini is
31    inference/reference only). Commercial-clean.
32
33    ## Ingested golden corpus ? `output/human_lovo_manifest.json` (n = 15)
34
35    | # | name | sport | rallies | fps | source proxy | video_id | labeled |
36    |---|---|---|---|---|---|---|---|
37    | 1 | adarsh_avi_singles | badminton (S) | 96 | 59.94 | (original 6) | ? | pre-06-16 |
38    | 2 | kushagra_singles | badminton (S) | 32 | 59.94 | (original 6) | ? | pre-06-16 |
39    | 3 | largetest_doubles | badminton (D) | 49 | 29.97 | (original 6) | ? | pre-06-16 |
40    | 4 | gbaaddy | badminton | 48 | 59.94 | (original 6) | ? | pre-06-16 |
41    | 5 | testlarge_short | badminton | 6 | 29.97 | (original 6) | ? | pre-06-16 |
42    | 6 | mahadevpura_singles | badminton (S) | 31 | 29.97 | (original 6) | ? | pre-06-16 |
43    | 7 | mahadevpura_2 | badminton (multicourt) | 40 | 59.94 | `~/Roora Request - June
15/[Done] Video 3 ~/mahadevpura-2_proxy.mp4` | `038e0e0915a2` | 2026-06-16 |
44    | 8 | GX010128 | badminton (multicourt) | 52 | 59.94 | `~/Golden Label Videos Request -
June 15/[Done] Video 13 ~/GX010128.MP4` | `db84f3e65318` | 2026-06-16 |
45    | 9 | mahadevpura_1 | badminton (S, continuous) | 10 | 59.94 | `~/[Done] Video 14 -
Badminton/mahadevpura-1.MP4` | `38bef75d78cb` | 2026-06-17 |
46
47    ## Labeled, ingestion in progress (WASB trajectory extracting)
48
49    | name | sport | rallies | fps | source proxy | status |
50    |---|---|---|---|---|---|
51    | GX030094 | badminton | 41 | **29.97** | `~/Golden Label ? /[Done] Video 6 -

```

```

Badminton/GX030094_proxy.mp4` | **ingested** (in `human_lovo_manifest.json`) |
52 | Badminton BXH 2 | badminton (multicourt) | 44 | 59.94 | `?/[Done] Video 10 - Badminton
Multicourt/Badminton BXH 2_proxy.mp4` | **ingested** (in `human_lovo_manifest.json`) |
53
54 ## Source folder & not-yet-labeled
55 - **Annotation root:** `G:\My Drive\Sharing\Annotation Videos\Golden Label Videos
Request - June 15\`
56 (+ the earlier `Roora Request - June 15\`). `[Done]` prefix in a dir name = golden-
labeled.
57 - **Labeled badminton dirs:** Video 3 (=mahadevpura-2), 6 (=GX030094), 10 (=BXH 2), 13
(=GX010128, a
58 re-label of #8), 14 (=mahadevpura-1). **Video 13 is a duplicate of #8** (already
ingested).
59 - **Not yet labeled / other sports:** Video 1, 4, 7, 9, 12 (badminton, pending); Video
2, 8 (table tennis);
60 Video 5, 11 (pickleball). Badminton is the current focus (alpha customers);
TT/pickleball labeled later.
61
62 ## How to isolate the golden videos
63 Each row's **source proxy** column is the file to copy into a dedicated golden-videos
location. The labels
64 live beside the proxy (`<name>.rallies.csv`) and a durable copy is in
`output/<name>.rallies.csv`; the
65 trajectory is in `Annotation Setup/Trajectories/` (older manifests may still say
`Collect/Trajectories/`;
66 use the path smoke checker above). Keep this table in sync on every new ingest.
67

```

12. user (2026-07-06T20:14:13.985Z)

```

1 # Phase-1 Golden-Set Video Collection Spec
2
3 **Status:** Design v1 . **Date:** 2026-06-02
4 **Companion docs:** [GOLDEN_SET_VENDORING.md](../archives/roora-
vendor/GOLDEN_SET_VENDORING.md) (vendor process, ADR-006) .
[VIDEO_RECORDING_GUIDELINES.md](VIDEO_RECORDING_GUIDELINES.md) (capture spec the user-facing
guide derives from) . scoring in [EVALUATION.md](../EVALUATION.md).
5
6 What footage to gather for the **first round** of golden-set creation so the trained
rally segmenter reaches good precision/recall ? and so the eval number actually predicts
production. Golden-set creation is **iterative**: expect multiple rounds, expanding toward
whatever axis round-1 metrics show is weakest.
7
8 ---
9
10 ## 1. The principle that dominates everything: match the deployment distribution
11
12 The CV literature is consistent: **models trained on broadcast footage fail to
generalize to self-recorded amateur video** ? broadcast uses a "straight, fixed camera with
minimum variability," while real non-elite footage has limited coverage, occlusions, and
inconsistent viewpoints ([team-sports
survey](https://link.springer.com/article/10.1007/s10462-024-10934-9), [unstructured-video event
detection](https://arxiv.org/pdf/2002.08097)). This is the **same failure mode as ADR-004**
(broadcast cameras *track players*, so player pixel-motion inverts during rallies).
13
14 Our product ingests **users' own match recordings** (a fixed phone/GoPro on a tripod).
Therefore:
15
16 - The golden set ? **especially the held-out eval split** ? must be dominated by
**static, self-recorded, full-court** footage that mirrors real users.
17 - **ShuttleSet / broadcast** footage is fine as *training* augmentation and as a
contrast set, but must be **minimal in eval**. Measuring precision/recall on broadcast yields a
number that does **not** predict production.
18
19 ---

```

```

20
21     ## 2. Primary recording profile (also the basis of the user-facing capture guide)
22
23     Grounded in the badminton-CV setup literature (camera ~2 m behind the baseline, ~4.5 m
elevated, ~30° down-tilt, whole court + headroom above the net ?
[source](https://pmc.ncbi.nlm.nih.gov/articles/PMC10747833/)):
24
25     | Setting | Spec | Why |
26     |---|---|---|
27     | Mount | Fixed tripod; **no pan/zoom** during play | Avoids the ADR-004 camera-tracking
trap |
28     | Position | Behind one baseline/end, centered | Full-court, symmetric view |
29     | Height | **375 m** elevated (~4.5 m ideal) | De-occludes the far player & net |
30     | Tilt | **~30° down** | Optimal coverage, low keystoneing |
31     | Framing | Entire court **+ headroom above the net**; players never leave frame |
Detection + court-relevance need it |
32     | Lighting | Even; no backlight/glare | Reduces blur & false motion |
33
34     Most of round 1 should use this profile; a minority deliberately varies it (§4) to probe
robustness.
35
36     ---
37
38     ## 3. Resolution & frame rate (per sport)
39
40     A controlled study: table-tennis ball detection was **22.5% @ 30 fps ? 93.2% @ 60 fps**
(120 fps ? 60 fps) ([Nature Sci Rep](https://www.nature.com/articles/s41598-023-42966-6));
OpenTTGames records at 120 fps. The ball/shuttle is the camera-invariant rally signal (our WASB
substrate, ADR-001), so fps is **not** a nicety for fast sports.
41
42     | Sport | Min resolution | Recommended fps | Note |
43     |---|---|---|---|
44     | Badminton | 1080p (720p floor) | 30 ok, **50/60 better** | Shuttle decelerates fast;
60 sharpens smashes |
45     | Pickleball | 1080p | **30?60** | Ball slower than TT; 30 likely fine, 60 safer |
46     | Table tennis | 1080p | **60 min (120 ideal)** | 30 fps collapses ball detection |
47
48     > These are **literature-backed starting points**; the published minimums get finalized
empirically against our golden set (VIDEO_RECORDING_GUIDELINES.md §3).
49
50     ---
51
52     ## 4. The Phase-1 diversity matrix
53
54     Cover the **high-impact axes** now; don't chase rare conditions in round 1. "Guards" =
whether the axis mostly protects **P**recision (not calling dead-time a rally) or **R**ecall
(not missing rallies).
55
56     | Dimension | Vary across? | Guards | Round-1 target |
57     |---|---|---|---|
58     | **Sport** | badminton, pickleball, table tennis | both | all 3 |
59     | **Format** | singles **and** doubles (badminton, pickleball) | R | both for those 2 |
60     | **Camera profile** | mostly primary-static; 1?2 side-elevated; 1?2 lower/phone-height;
**1?2 broadcast (contrast only)** | both | ~80% static / ~20% variants |
61     | **Venue / court color** | wood, green, blue synthetic | P | ?3 venues |
62     | **Lighting** | bright gym, dim club, mixed/backlit | both | ?3 conditions |
63     | **Background clutter** | empty hall vs people moving behind court | **P** (crowd
motion ? false positives) | both |
64     | **Players** | skill, attire colors, body types, L/R-handed, skin tone/age | R |
spread; don't reuse one pair |
65     | **Rally edge cases** | service faults, lets/replays, 1?2-shot & long rallies, net-
cords, coaching pauses, warm-up | boundary **P/R** | ensure present |
66     | **Resolution/fps ablation** | same match re-encoded 480/720/1080p; 30 vs 60 fps |
calibration | 1 matched set/sport |
67

```

```

68     ---
69
70     ## 5. Split it so the precision/recall number is honest (no train/eval leakage)
71
72     The #1 video-ML mistake: the same match in train and eval inflates the score ([data-
leakage guide](https://towardsdatascience.com/data-science-mistakes-to-avoid-data-
leakage-e447f88aae1c/)).
73
74     - Split by match / recording session ? never by clip. All clips from one match go
entirely to train or eval.
75     - Hold out whole conditions (a venue, a camera profile, a player pair the model
never saw) so eval measures generalization, not memorization.
76     - Reserve ~? of each sport as a deployment-like held-out eval set (static, self-
recorded). Broadcast/ShuttleSet may enter train but should be minimal in eval.
77
78     > Guardrails ? TODO (design with the user): train and eval live in separate
catalogs today, kept apart by convention. We want hard guardrails ? permissions / explicit
approval / re-confirmation at ingest & export ? so a video can never silently land in both.
Tracked in BACKLOG.md; to be designed in a later session.
79
80     ---
81
82     ## 6. Concrete "round 1" shopping list (per sport)
83
84     Aligns with the 12?15-videos/sport scale plan; the 3-clip Roora calibration is a subset
of this.
85
86     - 8 static-primary clips spanning venues / lighting / skill (?4 singles, 4 doubles
where applicable)
87     - 2 camera-variant clips (side-elevated, lower height) ? robustness probe
88     - 1 resolution/fps ablation set (one match @ 480/720/1080p, 30 & 60 fps)
89     - 1 broadcast contrast clip (ShuttleSet) ? train/contrast only
90     - Each ~2?4 min, deliberately including the edge cases above ? ~12 clips/sport;
hold out 4 whole matches for eval.
91
92     ---
93
94     ## 7. Deliberately defer to later rounds
95
96     Outdoor courts, phone-handheld/jitter, extreme angles, 4K, unusual net/court variants,
crowd-heavy tournaments. Add these once round-1 precision/recall shows the weakest axis, then
expand toward it (find the quality "knee" ? VIDEO_RECORDING_GUIDELINES.md §3).
97
98     ---
99
100    ## 8. Two cautions specific to precision (the stated goal)
101
102    - Backgrounds with people moving behind the court are the biggest false-positive
driver. Include some so the model learns to reject them ? and make sure the annotation
guideline's "ignore all dead time" rule is airtight.
103    - Warm-up / knock-up / practice swings look like rallies. Capture some, and confirm
the vendor excludes them (the Roora guide now has an explicit rule + a how-to-spot-it tip).
Early precision is bounded more by guideline / inter-annotator consistency (IAA) than by raw
data volume.
104
105    ---
106
107    ## 9. Sources
108
109    - [Team-sports action-recognition survey ? Springer
2024](https://link.springer.com/article/10.1007/s10462-024-10934-9)
110    - [Unstructured-video event detection ? arXiv
2002.08097](https://arxiv.org/pdf/2002.08097)
111    - [Badminton service-height CV setup (camera 4.5 m / 30°) ?
PMC](https://pmc.ncbi.nlm.nih.gov/articles/PMC10747833/)

```

```

112 - [VideoBadminton dataset ? arXiv 2403.12385](https://arxiv.org/html/2403.12385v1)
113 - [Table-tennis frame-rate vs detection accuracy ? Nature Sci
Rep](https://www.nature.com/articles/s41598-023-42966-6)
114 - [OpenTTGames (120 fps) dataset](https://lab.osai.ai/)
115 - [Action-spotting / segmentation survey ? arXiv
2505.03991](https://arxiv.org/html/2505.03991v2)
116 - [Data-leakage in train/test splits](https://towardsdatascience.com/data-science-
mistakes-to-avoid-data-leakage-e447f88aae1c/)
117

```

13. user (2026-07-06T20:14:14.794Z)

```

1      # Golden-Set Subproject ? Design & Implementation Overview
2
3      **Status:** In progress · **Date:** 2026-05-30
4      **Companion docs:** design rationale + vendor shortlist in
[GOLDEN_SET_VENDORING.md](../archives/roora-vendor/GOLDEN_SET_VENDORING.md) (ADR-006); scoring
mechanics in [EVALUATION.md](../EVALUATION.md).
5
6      This is the **implementation tracker** for turning the golden-set/vendor *design* into
shipped code. It's split into small, independently-reviewable PRs. The whole subproject is
**vendor-agnostic by construction** ? only one slice (GS-3, `HumanVendorProvider`) depends on
which vendor you pick offline; everything else can land first.
7
8      > **Update (2026-06-02):** Vendor selected for the **initial pilot ? Roora**
(Bengaluru). The **annotation guideline, intake-manifest template, and NDA** are authored in the
collector repo `docs/annotation/`, and the collector's `import-local` ingestion now covers the
**3 pilot sports** (badminton, pickleball, table tennis). On the indexer side, GS-1
(`file_provider.py`), GS-2 (`backend/eval/regression.py` + `run_regression.py`), and the GS-4
oracle scaffold (`automated_provider.py`) have landed; the remaining slice is the vendor-
dependent **GS-3 `HumanVendorProvider`**, parked until the Roora pilot confirms quality.
9
10     ---
11
12     ## Architecture in one picture
13
14     ```
15         video_ref ??? AnnotationProvider.submit()/poll()/fetch() ???
List[(start,end)]   ground truth
16
17         ?           ?           ?           ?
18         FileProvider HumanVendor Automated(oracle) ?
19         (GS-1)       (GS-3)       (GS-4)
20         backend.eval.metrics.score() ??? existing, unchanged
21
22         production pipeline ??? DB ??? backend.eval.harness.get_predictions() ???
List[(start,end)]   predictions
23
24         ?
25         regress(video, tier) (GS-2)
26
27         compares the two
28         ```
29
30     **The keystone:** a provider's `fetch()` returns the **same `List[(start, end)]`** that
`eval.annotations.load_annotations()` already returns and `eval.metrics.score()` already
consumes. So one scorer grades all three of: product-vs-GT, vendor-vs-answer-key, and annotator-
vs-annotator (IAA). No new scoring code ? only new *sources* of intervals.
31
32     ---
33
34     ## PR breakdown
35
36     | Slice | Branch | Scope | Depends on | Vendor-dependent? |
37     |---|---|---|---|---|
38     | **GS-1** | `feat/annotation-provider-registry` | `AnnotationProvider` ABC + registry +
`FileProvider` + tests | ? | No |
39     | **GS-2** | `feat/regression-runner` | `regress(video, tier, tolerance)` over existing

```

```

harness + `baseline.json` snapshot/compare + CLI | GS-1 | No |
35 | **GS-3** | `feat/human-vendor-provider` | `HumanVendorProvider` (submit/poll/fetch
against the chosen vendor's API or file-drop), DPA/secure-link handling | GS-1, **your vendor
pick** | **Yes** |
36 | **GS-4** | `feat/automated-oracle-provider` | `AutomatedProvider` fast-tier oracle ?
**interface stub** (injected labeler + fake; concrete model = backlog) | GS-1, GS-2 | No |
37 | **GS-5** | `feat/regression-triggers` | Wire triggers: CI gate, on-demand CLI,
scheduled run | GS-2, GS-4 | No |
38
39 Each row is one PR. GS-1 ? GS-2 ? GS-4 ? GS-5 can proceed **now**, independent of vendor
selection. GS-3 is parked until you've chosen a vendor (then it's a thin adapter behind the GS-1
interface).
40
41 ---
42
43 ## Slice details
44
45 ### GS-1 ? Provider registry + FileProvider ? (this PR)
46 - `backend/annotations/base.py` ? `AnnotationProvider` ABC (`submit`/`poll`/`fetch` +
`annotate()` convenience) and `annotation_registry`, mirroring `backend/providers` and
`backend/indexer` registry patterns.
47 - `backend/annotations/file_provider.py` ? `FileProvider`: the "job" is a CSV that
already exists (formalizes today's manual/ShuttleSet flow); reuses
`eval.annotations.load_annotations`.
48 - Tests assert the registry contract **and** the keystone (provider output is directly
scorable by `metrics.score()`).
49 - **Design choice:** `fetch`/`annotate` raise loudly on a missing/incomplete job
rather than returning `[]`, so a broken label source can't silently masquerade as "0 rallies"
and skew a regression to green.
50
51 ### GS-2 ? Regression runner (next)
52 `regress(video, tier="fast", tolerance=0.05)`: get GT from
`annotation_registry.get(tier)`, read predictions via the existing `eval.harness`, score with
the existing `metrics`, compare against a per-video `baseline.json`, and flag a regression when
precision/recall drop beyond tolerance. Pure orchestration over existing parts + a small
baseline store. CLI: `python run_regression.py --video ... --tier file`.
53
54 ### GS-3 ? HumanVendorProvider (after your vendor pick)
55 A concrete provider that submits a (licensed/owned-only) video to the chosen vendor and
fetches the returned rally CSV. Carries the privacy controls from the design doc (secure time-
limited links, delete-on-delivery, DPA). Thin adapter ? the interface already exists from GS-1.
56
57 ### GS-4 ? AutomatedProvider (fast-tier oracle) ?? oracle rule
58 An independent strong labeler for the fast/CI tier. **Must be a different model lineage
than whatever production uses** (else shared blind spots ? false green). A fast-tier failure is
a *smoke alarm* that escalates to the GS-3 human tier, never an authoritative verdict on its
own.
59
60 ### GS-5 ? Regression triggers
61 Wire `regress()` into CI (pre-merge gate on pipeline changes), an on-demand CLI (the
"surprise regression"), and an optional scheduled run. Includes a deliberate re-snapshot path so
legitimate production improvements can rebaseline instead of raising false regressions.
62
63 ---
64
65 ## Conventions
66 - **Branch per slice**, one PR each, tests green before PR, full suite must stay green.
67 - New runtime deps: **none** through GS-2 (GS-3 may add a vendor SDK; GS-4 may reuse the
existing `google-genai` for a Gemini oracle).
68 - All ground-truth video stays **licensed/owned** ? see the privacy section of the
design doc.
69

```

I now have all six files read fully. Let me note one important discrepancy to verify: the label CSV format differs between docs. `EVALUATION.md` says `start,end`; `GOLDEN\_VIDEOS.md` says `rally\_number,start\_time,end\_time,ending\_reason,sport,shots\_count`; `DATA\_IN\_GCS.md` uses `

15. user (2026-07-06T20:14:19.788Z)

backend\eval\gt_loader.py

16. user (2026-07-06T20:14:19.993Z)

backend\eval\annotations.py

17. assistant (2026-07-06T20:14:22.492Z)

Both exist. Let me read them to resolve the CSV column discrepancy authoritatively.

18. user (2026-07-06T20:14:23.412Z)

```
1      """Canonical ground-truth interval loader ? THE one GT/golden-CSV reader (ADR-008
prereq).
2
3      The repo historically carried two incompatible loaders:
4      - ``annotations.load_annotations``: positional ``start,end`` (the collector's
5        ``curate export`` format), raises on bad rows;
6      - ``calibrate_wasb.load_golden_gts``: DictReader requiring ``start_time/end_time``
7        (the rally-annotator golden format), silently skipping bad rows ? fed 7 modules.
8
9      A file in one convention fed to the other loader yielded zero/garbage intervals with no
10     error. This module accepts BOTH conventions (header-mapped, else positional), so every
11     quantity the ship-gate depends on (gt_hash, LOVO scores, the future consent fingerprint)
12     flows through one reader. Both legacy entry points now delegate here.
13
14     Accepted shapes (one rally per row; blank lines and ``#`` comments ignored):
15     - header with ``start``/``end`` OR ``start_time``/``end_time`` columns (any other
16       columns ? rally_number, ending_reason, sport, ? ? are ignored);
17     - no header: columns 0,1 positionally.
18     Timestamps: decimal seconds or ``MM:SS`` / ``HH:MM:SS``
19     (``annotations.parse_timestamp``).
20
21     """
22
23     import csv
24     import logging
25     from typing import List, Optional, Tuple
26
27     from backend.eval.annotations import parse_timestamp
28
29     logger = logging.getLogger(__name__)
30
31     Interval = Tuple[float, float]
32
33     _START_NAMES = ("start", "start_time")
34     _END_NAMES = ("end", "end_time")
35
36     def _header_columns(row: List[str]) -> Optional[Tuple[int, int]]:
37         """If `row` is a header naming start/end columns, return their indices, else
38         None."""
39         cells = [(c or "").strip().lower() for c in row]
40         start_idx = next((i for i, c in enumerate(cells) if c in _START_NAMES), None)
41         end_idx = next((i for i, c in enumerate(cells) if c in _END_NAMES), None)
42         if start_idx is not None and end_idx is not None:
43             return start_idx, end_idx
```

```

42         return None
43
44
45     def load_gt_intervals(csv_path: str, *, strict: bool = True) -> List[Interval]:
46         """Load GT rally intervals from either golden-CSV convention.
47
48         strict=True (corpus/gate use): raise ValueError with file:line on any malformed row
49         (bad timestamp, end <= start, negative start) so labeling mistakes fail loudly.
50         strict=False (legacy tuning-driver behavior): skip bad rows, but LOG how many were
51         skipped ? a silently shrinking GT set flatters precision.
52
53         Returns intervals sorted by start; overlapping GT intervals are logged as a warning
54         (they break one-to-one matching and usually indicate a labeling error).
55         """
56         intervals: List[Interval] = []
57         skipped = 0
58         cols: Optional[Tuple[int, int]] = None # (start_idx, end_idx); None until detected
59         saw_header = False
60
61         with open(csv_path, newline="", encoding="utf-8-sig") as f:
62             for lineno, row in enumerate(csv.reader(f), start=1):
63                 if not row or not (row[0] or "").strip() or row[0].strip().startswith("#"):
64                     continue
65                 if cols is None and not saw_header:
66                     header = _header_columns(row)
67                     if header is not None:
68                         cols = header
69                         saw_header = True
70                         continue # consume the header row
71                 cols = (0, 1) # headerless: positional start,end
72                 assert cols is not None # every branch above assigns it
73                 si, ei = cols
74                 try:
75                     if len(row) <= max(si, ei):
76                         raise ValueError(
77                             f"expected at least {max(si, ei) + 1} columns, got {row!r}"
78                         )
79                     start = parse_timestamp(row[si])
80                     end = parse_timestamp(row[ei])
81                     if start < 0:
82                         raise ValueError(f"negative start ({start})")
83                     if start >= end:
84                         raise ValueError(f"start ({start}) >= end ({end})")
85                 except ValueError as e:
86                     if strict:
87                         raise ValueError(f"{csv_path}:{lineno}: {e}") from e
88                     skipped += 1
89                     continue
90                 intervals.append((start, end))
91
92         if skipped:
93             logger.warning(
94                 "%s: skipped %d malformed GT row(s) (strict=False)", csv_path, skipped
95             )
96
97         intervals.sort(key=lambda iv: iv[0])
98         for (s1, e1), (s2, e2) in zip(intervals, intervals[1:]):
99             if s2 < e1:
100                 logger.warning(
101                     "%s: overlapping GT intervals (%.3f-%.3f vs %.3f-%.3f) ? "
102                     "labeling error? Overlaps break one-to-one matching.",
103                     csv_path,
104                     s1,
105                     e1,
106                     s2,

```



```

107             e2,
108         )
109     return intervals
110

```

19. assistant (2026-07-06T20:15:48.191Z)

I now have everything needed to resolve the CSV format discrepancy authoritatively. I have all the facts. Here is the consolidated report.

Report: Golden eval videos, labels, quick-run clips, and F1/tolF1 definitions

Five docs read in full plus the authoritative loader `backend/eval/gt\_loader.py`. Note upfront: this repo describes **two parallel worlds** ? (1) a legacy **local-file-system / SQLite-DB / OneDrive-Drive** world (the older docs), and (2) a newer **GCS bucket** world (`gs://khehsutra-rally-corpus`) that is mid-migration. I flag which is which throughout.

(a) Golden-video locations, filenames, and GCS URIs

GCS bucket (the current source-of-truth target)

Evidence: `docs\DATA\_IN\_GCS.md`

- **Bucket:** `gs://khehsutra-rally-corpus` (single corpus bucket), project `gen-lang-client-0306026826`, region `asia-south1`. ? lines 30-33.
- **Canonical layout contract** ? lines 48-55:
 - `videos/` holds `.mp4` ? the video prefix.
 - `labels/` holds `.rallies.csv` ? golden rally labels.
 - `trajectories/` holds cached WASB trajectory CSVs (the GPU-free eval input, consumed via `manifests/golden.json`).
 - `manifests/golden.json` ? the eval manifest (15 clips, `gs://` trajectory + label paths) ? lines 207-209.
- **Example URIs** (from the worker command block, lines 109-117): `gs://khehsutra-rally-corpus/videos/<clip>.mp4`, out-URI `gs://khehsutra-rally-corpus`.
- **Reality vs. contract:** only the **metadata** (15 trajectories + 15 canonical `labels/<date><name>.rallies.csv` + `manifests/golden.json`) actually landed in the bucket (Phase-1, 2026-06-22). The **videos** are almost entirely NOT migrated ? **"1/15"** demo'd (`mahadevpura\_1`); 14 remain" ? lines 275. As of the pre-Phase-1 snapshot, `videos/` held only 3 objects: `GX030094\_proxy.mp4`, `mahadevpura-1.MP4`, `mahadevpura\_smoke\_180s.mp4` ? line 77.
- **Owned-source video MD5 registry** (durable manifest, not the videos themselves): `sports-data-collector/data/golden\_source\_videos.json` (six owned golden MP4s) ? lines 139, 218-220.

Local / Drive world (where the actual golden video proxies live today)

Evidence: `docs\data\_pipeline\GOLDEN\_VIDEOS.md`

- **Annotation root:** `G:\My Drive\Sharing\Annotation Videos\Golden Label Videos Request - June 15\` (+ earlier `Roora Request - June 15\`). `[Done]` prefix = golden-labeled. ? line 55.
- **Actual per-clip source-proxy paths** (the "source proxy" column, lines 43-45, 51-52):
 - `mahadevpura\_2` ? `?/Roora Request - June 15/[Done] Video 3 ?/mahadevpura-2\_proxy.mp4`, video_id `038e0e0915a2`
 - `GX010128` ? `?/Golden Label Videos Request - June 15/[Done] Video 13 ?/GX010128.MP4`, video_id `db84f3e65318`
 - `mahadevpura\_1` ? `?/[Done] Video 14 - Badminton/mahadevpura-1.MP4`, video_id `38bef75d78cb`
 - `GX030094` ? `?/[Done] Video 6 - Badminton/GX030094\_proxy.mp4`
 - `Badminton BXH 2` ? `?/[Done] Video 10 - Badminton Multicourt/Badminton BXH 2\_proxy.mp4`
- **Phase-2 migration source table** (`docs\DATA\_IN\_GCS.md` lines 168-170): 7 Drive proxies (`mahadevpura\_1/2`, `GX010128`, `GX020094`, `GX030094`, `Badminton\_BXH\_2`, `Boxhill\_Doubles`, ~13.5 GB) + 8 F:-archive full MP4s (`GX010137`, `GX010141`, + original-6: `adarsh\_avi\_singles`, `kushagra\_singles`, `largetest\_doubles`, `gbaaddy`, `testlarge\_short`, `mahadevpura\_singles`, ~50 GB). F: root: `F:\[Khehsutra] GoPro Hero Black 12\GoPro Recorded Videos`.

De-attributed regression fixtures (a third, video-free artifact ? not videos)

Evidence: `docs\GOLDEN\_REGRESSION\_FIXTURES.md`

- Tier-0 fixtures live in-repo at `eval\_baselines/fixtures/` (tracked) ? line 87. Six real fixtures `fx\_r01`?`fx\_r06` (~1.9 MB total), de-attributed per-frame `trajectory.csv`, both 29.97 & 59.94 fps ? lines 94, 159, 174. These are **not** playable videos; they are frozen numeric trajectory streams for cross-machine regression. The surrogate?source name map is kept **off-git** (owner-held).

(b) Ground-truth labels: location, CSV format, and video?label pairing

Location

- **GCS:** `labels/<YYYY-MM-DD>\_<name>.rallies.csv` ? `docs\DATA\_IN\_GCS.md` line 199 (naming LOCKED decision; `<YYYY-MM-DD>` = label authoring date = earliest mtime).
- **Local durable copies:** `output/<name>.rallies.csv`, plus beside the proxy as `<name>.rallies.csv` ? `GOLDEN\_VIDEOS.md` lines 64-66.
- **Manifest tying it together:** `output/human\_lovo\_manifest.json` (local, n=15) and `gs://?/manifests/golden.json` (GCS, 15 clips).

CSV format ? ? TWO conventions, both accepted (this is a real, load-bearing discrepancy)

The docs disagree, and I resolved it against the authoritative loader

`backend\eval\gt\_loader.py`:

- **EVALUATION.md** (lines 27-31): **positional/headered** `start,end` ? decimal seconds or `MM:SS`/`HH:MM:SS`, header optional/auto-detected. This is the collector's `curate export` format.
- **GOLDEN_VIDEOS.md** (lines 15-16): **rally-annotator CSV** with columns `rally\_number,start\_time,end\_time,ending\_reason,sport,shots\_count` (decimal seconds), ingested by `backend/eval/gt\_loader.py` which reads `start\_time`/`end\_time`.
- **Authoritative resolution** ? `gt\_loader.py` accepts BOTH (it was written specifically to unify two historically-incompatible loaders ? docstring lines 3-18). Rules (lines 31-32, 45-90): it matches a header naming `start`/`start\_time` (start col) and `end`/`end\_time` (end col); **all other columns** ? `rally\_number`, `ending\_reason`, `sport`, `shots\_count` ? are ignored; if no header, uses columns 0,1 positionally. `strict=True` (corpus/gate) raises `file:line` on any bad row; `strict=False` (legacy) skips+logs. So the two "different" formats are the same start/end pair with extra ignored columns.

Video ? label pairing convention

- **Filename convention:** `<name>.rallies.csv` pairs with the clip named `<name>` ? confirmed in `DATA\_IN\_GCS.md` layout (`videos/<clip>.mp4` ? `labels/<clip>.rallies.csv`, lines 49-50) and `GOLDEN\_VIDEOS.md` line 64. GCS labels additionally carry a date prefix;
`backend/worker/\_\_main\_\_.py::label\_clip\_name()` strips both the `.rallies.csv`/`.csv` suffix **and** a leading `^d{4}-d{2}-d{2}\_` prefix so `2026-06-22\_GX020094.rallies.csv` resolves to bare `videos/GX020094` ? `DATA\_IN\_GCS.md` lines 238-246.
- **Strong keying (MD5):** `video\_id` = MD5 of the proxy via `backend/utils/hashing.py::compute\_video\_id`; it keys the WASB trajectory cache, Gemini reference rows, and golden features ? `GOLDEN\_VIDEOS.md` lines 17-19. The eval CLI keys the DB by file MD5 (`--video <file>` computes it, or `--video-id <md5>`) ? `EVALUATION.md` lines 51-57.
- **Alignment caveat:** label timebase must match the windowed proxy (same fps + frame count); verify with `ffprobe` ? the `largetest\_doubles` "near-miss" is cited as why ? `GOLDEN\_VIDEOS.md` lines 18-19.

(c) 2-3 specific clips for a quick numbers run (with resolution / fps)

Best candidates, drawn from the ingested corpus table in `GOLDEN\_VIDEOS.md` (lines 37-52) and the R2 preview clips:

Clip	Sport / type	Rallies	fps	Notes / evidence
---	---	---	---	---
GX010128	badminton, multicourt	52	59.94	Has a fully worked tolF1 preview (see (d)); video_id `db84f3e65318`. `GOLDEN_VIDEOS.md:44`; `R2_EVAL_METRIC_DESIGN.md:108-109`.
mahadevpura_2	badminton, multicourt	40	59.94	The primary n=2 preview clip;

```
video_id `038e0e0915a2`. `GOLDEN_VIDEOS.md:43`; `R2_EVAL_METRIC_DESIGN.md:106-107`. |
| mahadevpura_1 | badminton (S, continuously-active blob) | 10 | 59.94 | Small (10
rallies), fast; reported as its own stratum. It's also the only clip already in the GCS
`videos/` bucket. video_id `38bef75d78cb`. `GOLDEN_VIDEOS.md:45`; `DATA_IN_GCS.md:275`. |
| (alt, smallest) testlarge_short | badminton | 6 | 29.97 | Fewest rallies of all; but
"original 6" with no video_id/GCS proxy recorded. `GOLDEN_VIDEOS.md:41`. |
```

Resolution: the docs give an explicit numeric fps per clip but **do not state per-clip pixel resolution** in these registries. The only resolution guidance is the spec target ? Badminton **1080p** (720p floor), 30 ok / 50-60 better ? `GOLDEN\_SET\_PHASE1\_VIDEOS.md` lines 44-46; the Drive proxies are noted as "already 1080p" (`DATA\_IN\_GCS.md` line 180). **Per-clip actual resolution is an open item** (see flags below). The named 3-clip "Roora calibration" subset is mentioned but not enumerated by name ? `GOLDEN\_SET\_PHASE1\_VIDEOS.md` line 84.

`n2` ground-truth clips used across the metric design = **mahadevpura-2** and **GX010128** (`R2\_EVAL\_METRIC\_DESIGN.md` line 34, table lines 106-109) ? the cleanest pair for a quick numbers run.

(d) F1 / tolF1 definition and the ? tolerance values

Two distinct metrics coexist. Do not conflate them.

1. tIoU-based F1 (the current gate) ? `docs\EVALUATION.md`

- Matching:** predictions and GT are `[start, end)` intervals; a prediction matches a GT rally when their **temporal IoU** (overlap ÷ union) ? **threshold**. Matching is **greedy, one-to-one** (highest-IoU pairs claimed first) ? lines 112-118, 5-line "How matching works" section.
- IoU threshold ?:** default **0.5** (`--iou`, default `0.5`); optional PR sweep 0.1-0.9 ? lines 56, 65-67.
- P / R / F1:** precision = TP/(TP+FP), recall = TP/(TP+FN), **F1** = their harmonic mean ? lines 87-89. Also reports mIoU (mean temporal IoU over matched pairs) and dStart/dEnd (mean absolute boundary error in seconds).
- Level:** rally-level (per-video `[start,end)` rally intervals). Scores predictions already stored in the SQLite DB; deterministic, no API calls ? lines 3-7.
- Single-video command** (`EVALUATION.md` lines 51-57):
 - `python run\_eval.py --video "C:/videos/match.mp4" --annotations rallies.csv`
 - or `python run\_eval.py --video-id <md5> --annotations rallies.csv --stage both --iou 0.5`
 - Inputs needed:** the video must already be processed so detections are in the DB (harness does **not** run the pipeline), plus the ground-truth `rallies.csv`. Keyed by file MD5.

2. tolF1 / F1@? (the newer "R2" tolerance metric) ? `docs\R2\_EVAL\_METRIC\_DESIGN.md`

- Definition** (the exact difference from tIoU-F1): instead of an IoU-overlap gate, a predicted window matches a golden rally **iff** `|?start| ? ?\_start` AND `|?end| ? ?\_end` (asymmetric boundary tolerance) ? line 52. Assignment is **strict 1:1 Hungarian** (`scipy.optimize.linear\_sum\_assignment`), **not** greedy/overlap ? lines 62-64. Then `P@?` = matched / n_pred, `R@?` = matched / n_ref, `F1@?` = harmonic mean ? line 69. Reported as a **decomposition** (P@?, R@?, F1@?, split start/end boundary-error distributions, over-seg guard, ?-sweep), never one number ? lines 68-88.
- values ? in SECONDS** (temporal, not pixels):
 - Provisional locked pair: **?_start = 2.0 s**, **?_end = 1.5 s** (asymmetric ? forgive the fuzzy start, keep the rally-end honest) ? lines 53-56, 95-96, 183-186.
 - ?-sweep set (symmetric, for the trend line): **? ? {1.5, 2.0, 2.5, 3.0}** seconds ? line 84.
 - These ? are explicitly **provisional pending an IAA** (inter-annotator agreement) study; final `?\_start`/`?\_end` should be set to ~the 90th-percentile pairwise human `|?start|`/`|?end|` ? §2.4, lines 89-96.
- tolF1 vs F1 relationship:** `tolF1@~3.0 ? tIoU-F1@0.5` ? the tolerance sweep "contains tIoU as a point" (continuity with the frozen baseline) ? lines 105, 112-113.
- Status / which is the gate:** R2 is **IMPLEMENTED** but default-additive ?
- backend/eval/tolerance_metrics.py** landed and `rally\_seg\_eval` reports the R2 block **alongside** tIoU; **tIoU@0.5** remains the gate pending an ablation experiment ? lines 2-6, 147-149, 187-195. So today the authoritative headline number is still tIoU-F1@0.5; tolF1 is reported next to it.

- CLI flags: `--tau-start` / `--tau-end`` landed on `rally_seg_eval``; the `?-sweep` helper `tolerance_metrics.tolerance_sweep`` exists but is not yet surfaced as a CLI flag ? lines 148, 167.

Open questions / things NOT found (flagged)

1. **Per-clip pixel resolution is not recorded** in any of these docs. The registries give fps per clip but only spec-level target resolution (1080p badminton). `DATA_IN_GCS.md`` line 191 even warns that `fps/frame_width` must be **re-probed** from the bucket videos in Phase-3 ("don't trust the copied values"). If exact resolution per clip is needed, it must be `ffprobe``d from the actual proxy ? it's not documented.
2. **The named 3-clip "Roora calibration" subset** is referenced (`GOLDEN_SET_PHASE1_VIDEOS.md:84``) but the three specific clip names are never enumerated. My recommended quick-run trio (GX010128, mahadevpura_2, mahadevpura_1) is inferred from the ingested-corpus table + the R2 preview, not from an explicit "quick set" list.
3. **CSV-format inconsistency across docs is real but harmless:** `EVALUATION.md`` says `start,end``; `GOLDEN_VIDEOS.md`` says the 6-column `rally_number,start_time,end_time,ending_reason,sport,shots_count``. `backend/eval/gt_loader.py`` reconciles both (reads start/end under either header, ignores extra columns). Cite `gt_loader.py:1-18,31-90`` as authoritative if this matters downstream.
4. **Video availability gap:** despite the GCS "source-of-truth" framing, only 1 of 15 golden videos is actually in `gs://?/videos/`` today. For a live numbers run you currently need the local Drive proxies (`G:\My Drive\...``) or the F: archive, plus a processed DB ? the GCS eval path (Phase-3) can't fetch `gs://`` yet (`DATA_IN_GCS.md:186-194``, blocker §7b#2).
5. **No `.MP4`` ? `.rallies.csv`` literal example** is given verbatim, but the convention is unambiguous from the layout table: `videos/<name>.mp4`` pairs with `labels/<name>.rallies.csv`` (GCS adds a `<date>_`` prefix on the label, stripped at join time).

All paths above are quoted from the docs; file:line regions are cited inline for each fact.