

claude-windows / 46eabfd2-0e71-4ff5-a68d-c76f474d4ff3

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3.jsonl`  
- SHA-256: `8dee64ffd48b9fe577a978cbc787c9cb07d7322b026c28fd5e4b704cbb9b51cf`  
- Source modified: `2026-07-08T09:46:16+00:00`  
- Imported at: `2026-07-08T15:59:19+00:00`  
- project: `C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer`  
- session_id: `46eabfd2-0e71-4ff5-a68d-c76f474d4ff3`
```

Transcript

1. user (2026-07-08T05:53:55.965Z)

Follow-up to #521 (PR #525, merged compile-on-L4 + NVENC). Part 3 of #521 was deferred: make the branding watermark CHEAPER so the per-clip stitch isn't a full 4K re-encode.

Context: `backend/pipeline/stitcher/compiler.py::VideoCompiler` burns the watermark via a `drawtext`/`movie` `-vf` filtergraph during a per-clip re-encode. #521 already added an `encoder` knob (NVENC on the L4, libx264 fallback) via `\_select\_encoder`, which hardware-accelerates the ENCODE ? but every clip is still fully re-encoded because the watermark filter forces it. The 2026-07-08 CUJ (\$9 of docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md) showed the watermark re-encode was the dominant stitch cost.

Task: investigate a path where the stitch is NOT a full re-encode:

- Option A: pre-render the watermark to a transparent PNG overlay ONCE (per resolution) and `overlay` it, still cheaper than drawtext-per-frame; or
- Option B: apply the watermark only at the CONCAT stage (single pass over the final reel) instead of per-clip, so per-clip trims can stream-copy (`-c copy`) ? the big win, since trims dominate. This changes the two-stage codec contract in `compile\_highlights` (per-clip trims currently re-encode so the concat can `-c copy`); flipping it means trims `-c copy` and the concat re-encodes once with the watermark. Watch: `-ss`/`-to` stream-copy cuts at keyframes (frame-accuracy loss) ? may need a hybrid (re-encode only the boundary GOPs) or keep the current approach for frame-accurate trims.
- Keep libx264 the default + parity-safe; measure on a real 4K clip. Respect D5/D16(b) parity notes in docs/COMPUTE_DECOUPLED_SERVING/README.md (reel-encode parity is already relaxed under NVENC opt-in).

Document the chosen approach + measured before/after in docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md \$3 (which already flags this as a tracked follow-up). Branch off origin/master; one PR referencing #521.

2. assistant (2026-07-08T05:54:04.426Z)

I'll start by ramping up on the current state. Let me pull the repo and read the key files in parallel, following the repo-freshness and handoff conventions.

3. user (2026-07-08T05:54:09.106Z)

On branch codex/issue-523-sd-reel-plan
Your branch is up to date with 'origin/master'.

Changes not staged for commit:

```
(use "git add <file>..." to update what will be committed)  
(use "git restore <file>..." to discard changes in working directory)  
    modified:   SESSION_HANDOFF.md
```

Untracked files:

```
(use "git add <file>..." to include in what will be committed)
score_golden_gate.py
score_local.py
```

no changes added to commit (use "git add" and/or "git commit -a")

---BRANCH---

codex/issue-523-sd-reel-plan

---FETCH---

4. user (2026-07-08T05:54:09.149Z)

```
1      # Session Handoff
2
3      ## You are here
4
5      This checkout is on `master` at `bf97700` after a successful `git pull --ff-only`
6      on 2026-07-08. The top-level workspace folder is not a Git repo; the relevant
7      engine repo is `badminton-highlight-indexer/`. There are pre-existing untracked
8      local files `score_golden_gate.py` and `score_local.py`; do not stage or delete
9      them.
10
11     A planning pass for the optional cheap SD reel request is complete. Tracking bug:
12     https://github.com/Khelsutra/badminton-highlight-indexer/issues/523
13
14     ## Next steps / open threads
15
16     1. Engine PR for #523: add an opt-in compile rendition field (`original` default,
17       `sd` optional), persist it through `jobs.params`, and make `_execute_compile`
18       call `VideoCompiler.compile_highlights(...)` exactly once using the requested
19       profile.
20     2. SPA PR in sibling `khelsutra/web`: expose the SD option near Build, send the
21       compile rendition only when selected, and use a distinct `-sd.mp4` filename.
22     3. Bundle follow-up: rebuild/import the SPA snapshot into `backend/ui` after the
23       source SPA PR lands, if bundling remains separate.
24     4. Link each PR back to #523 and add verification notes/artifacts there.
25
26     ## Ramp-up kit
27
28     - `backend/main.py::CompileRequest` and `backend/main.py::compile_highlights`
29     - `backend/orchestration.py::_resolve_compile` and `::_execute_compile`
30     - `backend/infrastructure/database_jobs.py::create_compile_job`
31     - `backend/pipeline/stitcher/compiler.py::VideoCompiler`
32     - `tests/test_api_endpoints.py` compile tests
33     - `tests/test_compiler.py`
34     - Sibling repo `../khelsutra/web/src/api/client.ts`
35     - Sibling repo `../khelsutra/web/src/App.tsx`
36     - Sibling repo `../khelsutra/web/src/components/SelectionFooter.tsx`
37
38     ## Key decisions
39
40     - SD is a replacement compile mode, not an additional output: if selected, build
41       and deliver only the SD reel.
42     - Backward compatibility matters: old compile payloads should behave as original
43       quality.
44     - Use the existing async compile job shape; no DB migration is needed because
45       compile-specific params already live in `jobs.params`.
46     - Apply the SD profile before any per-rally ffmpeg trim so temporary clips are
47       consistently encoded and the final concat can remain stream-copy.
48     - Prefer typed, tunable stitching SD settings over hard-coded encode constants.
49
```

5. user (2026-07-08T05:54:09.167Z)

```
1      import logging
2      import os
```

```

3     import subprocess
4     import tempfile
5     from pathlib import Path
6     from typing import Any, Dict, List, Optional, Tuple, Union
7
8     from backend.utils.ffmpeg import ( # P2a: single ffmpeg/ffprobe resolver
9         ffmpeg_bin,
10        ffmpeg_timeout_sec,
11        ffprobe_bin,
12        ffprobe_timeout_sec,
13        timeout_label,
14    )
15
16    logger = logging.getLogger(__name__)
17
18    # Bundled, commercially-clean default watermark font (Apache-2.0). Shipped so the
overlay
19    # renders out-of-the-box without a system fontconfig setup (a common gap on Windows,
where
20    # `font='Sans` fails). Overridable via stitching.watermark.font_path. This module lives
at
21    # backend/pipeline/stitcher/, so parents[2] == backend/.
22    _BUNDLED_FONT = (
23        Path(__file__).resolve().parents[2] / "assets" / "fonts" / "RobotoSlab-Regular.ttf"
24    )
25
26
27    class VideoCompiler:
28        def __init__(
29            self,
30            output_dir: str,
31            end_padding: float = 1.0,
32            begin_padding: float = 0.5,
33            watermark: Optional[Any] = None,
34        ):
35            self.output_dir = output_dir
36            self.end_padding = end_padding
37            self.begin_padding = begin_padding
38            # Optional branding overlay. Accepts a WatermarkConfig model or a plain dict;
39            # normalized to a dict (or None) so this module stays config-package-agnostic.
40            self.watermark = self._normalize_watermark(watermark)
41            os.makedirs(output_dir, exist_ok=True)
42
43        @staticmethod
44        def _normalize_watermark(watermark: Optional[Any]) -> Optional[Dict[str, Any]]:
45            if watermark is None:
46                return None
47            if hasattr(watermark, "model_dump"):
48                return watermark.model_dump()
49            if isinstance(watermark, dict):
50                return dict(watermark)
51            return None
52
53        @staticmethod
54        def _ff_quote(path: str) -> str:
55            """Make a filesystem path safe to embed inside single quotes in an ffmpeg
56            filtergraph operand (drawtext fontfile/textfile, movie= source): forward
57            slashes (Windows-friendly), escaped single quotes, AND escaped colons.
58
59            The colon MUST be escaped even inside single quotes: ffmpeg's filtergraph
60            option parser treats ':' as the option separator, so a Windows drive-letter
61            path like ``C:/dir/font.ttf`` otherwise truncates the operand at ``C`` and
62            the encode fails with "Invalid argument". This silently disabled the
63            on-by-default bundled-font watermark (and absolute logo paths) on Windows."""
64            return path.replace("\\", "/").replace("'", r"\'").replace(":", r"\:")

```

```

65
66 @staticmethod
67 def _parse_time(val: Union[int, float, str]) -> float:
68     """Normalizes a timestamp value to float seconds.
69     Handles: float/int (pass-through), plain numeric strings ("12.4"),
70     and colon-separated strings ("MM:SS", "HH:MM:SS", "H:MM:SS.ms").
71     This mirrors the parse_time logic used in gemini.py to ensure
72     the stitcher can always consume whatever format the indexer produced."""
73     if isinstance(val, (int, float)):
74         return float(val)
75     if isinstance(val, str):
76         val = val.strip()
77         if ":" in val:
78             # Handle MM:SS, HH:MM:SS, etc.
79             parts = val.split(":")
80             parts.reverse()
81             total = 0.0
82             for i, p in enumerate(parts):
83                 try:
84                     total += float(p) * (60**i)
85                 except ValueError:
86                     pass
87             return total
88         try:
89             return float(val)
90         except ValueError:
91             logger.warning(
92                 f"Could not parse timestamp value '{val}' as a number. Defaulting to
93                 0.0."
94             )
95             return 0.0
96         logger.warning(
97             f"Unexpected timestamp type {type(val).__name__} for value '{val}'.
98             Defaulting to 0.0."
99         )
100         return 0.0
101
102 @staticmethod
103 def _merge_overlapping(
104     segments: List[Tuple[float, float]],
105 ) -> List[Tuple[float, float]]:
106     """Normalize, sort, and merge overlapping/duplicate (start, end) ranges.
107
108     A highlight reel must never replay the same footage. The same rally can land in
109     the segment list more than once ? e.g. two detector runs accumulated for one
110     video (add_video is an UPSERT, so re-running without a prune leaves both sets),
111     or
112     chunk-overlap duplicates near boundaries. Stitching those verbatim plays the
113     rally
114     twice. Merging any range that overlaps or abuts the previous one collapses true
115     duplicates/overlaps while preserving genuinely distinct (gapped) rallies.
116     Degenerate rows (end <= start) are dropped. Timestamps may be float/int/str.
117     """
118     parsed = []
119     for raw_s, raw_e in segments:
120         s = VideoCompiler._parse_time(raw_s)
121         e = VideoCompiler._parse_time(raw_e)
122         if e > s:
123             parsed.append((s, e))
124     parsed.sort()
125     merged: List[Tuple[float, float]] = []
126     for s, e in parsed:
127         if merged and s <= merged[-1][1]: # overlaps/abuts previous -> extend it
128             merged[-1] = (merged[-1][0], max(merged[-1][1], e))
129         else:

```

```

126         merged.append((s, e))
127     return merged
128
129 def _get_video_duration(self, video_path: str) -> float:
130     """Probes the video file to get its total duration in seconds.
131     Returns 0.0 if it cannot be determined."""
132     try:
133         timeout = ffprobe_timeout_sec()
134         result = subprocess.run(
135             [
136                 ffprobe_bin(),
137                 "-v",
138                 "error",
139                 "-show_entries",
140                 "format=duration",
141                 "-of",
142                 "default=noprint_wrappers=1:nokey=1",
143                 video_path,
144             ],
145             capture_output=True,
146             text=True,
147             check=True,
148             timeout=timeout,
149         )
150         return float(result.stdout.strip())
151     except subprocess.TimeoutExpired:
152         logger.warning(
153             "Could not determine video duration via ffprobe: timed out after %s",
154             timeout_label(ffprobe_timeout_sec()),
155         )
156         return 0.0
157     except Exception as e:
158         logger.warning(f"Could not determine video duration via ffprobe: {e}")
159         return 0.0
160
161 def _watermark_filter(self, tmp_dir: str) -> Optional[str]:
162     """Build the ffmpeg -vf filtergraph that burns the configured watermark (text
163     and/or logo PNG) into a clip during the per-clip re-encode. Returns None when
164     watermarking is disabled or has nothing to draw. The concat stage stays -c copy
165     because the mark is already baked into every clip identically."""
166     wm = self.watermark
167     if not wm or not wm.get("enabled"):
168         return None
169
170     text = (wm.get("text") or "").strip()
171     logo_path = (wm.get("logo_path") or "").strip()
172     if logo_path and not os.path.exists(logo_path):
173         logger.warning(
174             f"[watermark] logo not found: {logo_path} ? skipping logo overlay."
175         )
176         logo_path = ""
177     if not text and not logo_path:
178         return None
179
180     margin = int(wm.get("margin", 24))
181     opacity = float(wm.get("opacity", 0.85))
182     position = str(wm.get("position", "bottom-right"))
183     # drawtext uses lowercase w/h (main) + tw/th (text); overlay uses W/H (main) +
w/h (logo).
184     dt_xy = {
185         "bottom-right": f"x=w-tw-{margin}:y=h-th-{margin}",
186         "bottom-left": f"x={margin}:y=h-th-{margin}",
187         "top-right": f"x=w-tw-{margin}:y={margin}",
188         "top-left": f"x={margin}:y={margin}",
189     }.get(position, f"x=w-tw-{margin}:y=h-th-{margin}")

```

```

190         ov_xy = {
191             "bottom-right": f"W-w-{{margin}}:H-h-{{margin}}",
192             "bottom-left": f"{{margin}}:H-h-{{margin}}",
193             "top-right": f"W-w-{{margin}}:{{margin}}",
194             "top-left": f"{{margin}}:{{margin}}",
195         }.get(position, f"W-w-{{margin}}:H-h-{{margin}}")
196
197         drawtext = ""
198         if text:
199             ratio = float(wm.get("font_size_ratio") or 0.035)
200             divisor = max(1, round(1.0 / ratio))
201             # textfile= sidesteps the notorious drawtext text= escaping for arbitrary
strings.
202             tf_path = os.path.join(tmp_dir, "watermark.txt")
203             with open(tf_path, "w", encoding="utf-8") as fh:
204                 fh.write(text)
205             # Font precedence: explicit font_path ? bundled Apache-2.0 default (renders
206             # without fontconfig) ? fontconfig family (last-resort, environment-
dependent).
207             font_path = (wm.get("font_path") or "").strip()
208             if not font_path and _BUNDLED_FONT.exists():
209                 font_path = str(_BUNDLED_FONT)
210             if font_path:
211                 font_opt = f"fontfile='{self._ff_quote(font_path)}'"
212             else:
213                 font_opt = f"font='{wm.get('font', 'Sans')}'"
214             box = ""
215             if wm.get("box", True):
216                 box = f"box=1:boxcolor=black@{{min(opacity, 0.5)}:.2f}:boxborderw=10"
217             drawtext = (
218                 f"drawtext=textfile='{self._ff_quote(tf_path)}':{font_opt}"
219                 f"fontcolor=white@{{opacity:.2f}}:fontsize=h/{{divisor}}:{{dt_xy}}{box}"
220             )
221
222             if logo_path and drawtext:
223                 return
224             f"movie='{self._ff_quote(logo_path)}'[wm];[in][wm]overlay={ov_xy},{drawtext}"
225             if logo_path:
226                 return f"movie='{self._ff_quote(logo_path)}'[wm];[in][wm]overlay={ov_xy}"
227             return drawtext
228
229     def _trim_one_segment(
230         self,
231         idx: int,
232         raw_start: Union[int, float, str],
233         raw_end: Union[int, float, str],
234         segments: List[Tuple[float, float]],
235         video_duration: float,
236         tmp_dir: str,
237         raw_video_path: str,
238         watermark_filter: Optional[str],
239     ) -> Tuple[Optional[str], Optional[Tuple[int, float, float, str]]]:
240         """Trim a single segment into its own clip file (extracted verbatim from the
241         per-segment body of compile_highlights). Returns (clip_path, skip_record): on
242         success the clip path and None; on a skip (invalid/out-of-range/ffmpeg failure)
243         None and the (idx, start, end, reason) record to record in skipped_segments."""
244         # Normalize timestamps to float seconds ? handles float, int,
245         # plain numeric strings, and colon-separated formats like "MM:SS"
246         start = self._parse_time(raw_start)
247         end = self._parse_time(raw_end)
248         segment_label = (
249             f"Segment #{idx + 1}/{len(segments)} [{start:.2f}s - {end:.2f}s]"
250         )
251
252         # Validate the segment times

```

```

252         if start < 0:
253             logger.warning(
254                 f"{segment_label}: start_time is negative ({start:.2f}s). Clamping to
0."
255             )
256             start = 0.0
257
258         if start >= end:
259             logger.error(
260                 f"{segment_label}: start_time ({start:.2f}s) >= end_time ({end:.2f}s).
Skipping invalid segment."
261             )
262             return None, (idx, start, end, "start >= end")
263
264         # Check if segment extends beyond video duration
265         if video_duration > 0:
266             if start >= video_duration:
267                 logger.error(
268                     f"{segment_label}: start_time ({start:.2f}s) is beyond the video
duration ({video_duration:.2f}s). "
269                     f"This segment cannot be extracted ? the video ran out. Skipping."
270                 )
271                 return None, (idx, start, end, "start beyond video end")
272
273             if end > video_duration:
274                 original_end = end
275                 end = video_duration
276                 logger.warning(
277                     f"{segment_label}: end_time ({original_end:.2f}s) exceeds video
duration ({video_duration:.2f}s). "
278                     f"Clamping end_time to {video_duration:.2f}s. The segment may be
truncated."
279                 )
280
281             clip_path = os.path.join(tmp_dir, f"clip_{idx}.mp4")
282
283             # Precise sub-second trim using fast re-encoding to preserve stream headers
284             padded_start = max(0.0, start - self.begin_padding)
285             padded_end = end + self.end_padding
286
287             # Clamp padded_end to video duration if known
288             if video_duration > 0 and padded_end > video_duration:
289                 padded_end = video_duration
290                 logger.info(
291                     f"{segment_label}: Padded end ({end + self.end_padding:.2f}s) exceeds
video duration. "
292                     f"Clamping to {video_duration:.2f}s (end padding partially applied)."
293                 )
294
295             trim_duration = padded_end - padded_start
296             logger.info(
297                 f"Trimming {segment_label} -> [{padded_start:.3f}s - {padded_end:.3f}s]
(duration: {trim_duration:.2f}s)"
298             )
299
300             base_cmd = [
301                 ffmpeg_bin(),
302                 "-y",
303                 "-ss",
304                 str(round(padded_start, 3)),
305                 "-to",
306                 str(round(padded_end, 3)),
307                 "-i",
308                 raw_video_path,
309                 "-c:v",

```

```

310         "libx264",
311         "-preset",
312         "superfast",
313         "-crf",
314         "22",
315         "-c:a",
316         "aac",
317         "-b:a",
318         "128k",
319     ]
320     vf_args = ["-vf", watermark_filter] if watermark_filter else []
321
322     # Try with the watermark first; if that encode fails (e.g. a bad font/logo
323     # path), retry once WITHOUT it so a branding misconfig can never nuke the
324     # whole reel. Clips stay codec-identical either way, so concat -c copy holds.
325     attempts = [vf_args, []] if vf_args else [[]]
326     produced = False
327     last_exc: Optional[subprocess.CallledProcessError] = None
328     cmd = base_cmd + [clip_path]
329     timeout = ffmpeg_timeout_sec()
330     for attempt_vf in attempts:
331         cmd = base_cmd + attempt_vf + [clip_path]
332         try:
333             subprocess.run(
334                 cmd, capture_output=True, check=True, timeout=timeout
335             )
336         except subprocess.TimeoutExpired:
337             reason = f"ffmpeg timeout after {timeout_label(timeout)}"
338             logger.error(f"{segment_label}: FFmpeg trim timed out.")
339             print(f"  Command: {' '.join(cmd)}")
340             print(f"  {reason}")
341             return None, (idx, start, end, reason)
342         except subprocess.CallledProcessError as e:
343             last_exc = e
344             if attempt_vf and vf_args:
345                 _err = (
346                     e.stderr.decode(errors="replace").strip()[-300:]
347                     if getattr(e, "stderr", None)
348                     else ""
349                 )
350                 logger.warning(
351                     f"{segment_label}: watermark encode failed; retrying without it."
352                     f"ffmpeg: {_err or '(no stderr captured)'}"
353                 )
354                 continue
355             break
356     if os.path.exists(clip_path) and os.path.getsize(clip_path) > 0:
357         produced = True
358         if not attempt_vf and vf_args:
359             logger.warning(
360                 f"{segment_label}: encoded WITHOUT the watermark (the watermark
361                 filter failed)."
362             )
363             break
364     last_exc = None # ffmpeg succeeded but produced nothing
365     break
366
367     if produced:
368         return clip_path, None
369     elif last_exc is not None:
370         stderr_msg = last_exc.stderr.decode(errors="replace").strip()
371         logger.error(f"{segment_label}: FFmpeg trim failed.")
372         print(f"  Command: {' '.join(cmd)}")
373         print(

```



```

373         f" FFmpeg stderr: {stderr_msg[-500:]]}"
374     ) # Last 500 chars to avoid flooding
375     return None, (idx, start, end, f"ffmpeg error: {stderr_msg[-200:]]}")
376 else:
377     logger.error(
378         f"{segment_label}: FFmpeg exited successfully but output clip is missing
or empty. Skipping."
379     )
380     return None, (idx, start, end, "empty output clip")
381
382 def compile_highlights(
383     self,
384     raw_video_path: str,
385     segments: List[Tuple[float, float]],
386     output_filename: str,
387 ) -> str:
388     """
389     Extracts selected segments from a raw video and stitches them together
390     into a seamless highlights file, applying end_padding to prevent abrupt endings.
391
392     .. warning::
393         This method uses a two-stage codec contract: per-clip re-encodes must use
394         IDENTICAL encoder settings (codec, preset, CRF, pixel format) so the final
395         concat can stream-copy (``-c copy``) without re-encoding. Changing the
396         clip encode settings without updating the concat assumptions will produce
397         broken output.
398     """
399     if not segments:
400         raise ValueError("No highlight segments provided to compile.")
401
402     # Never stitch the same footage twice: collapse overlapping/duplicate ranges
403     # (e.g. accumulated re-run segments or chunk-overlap dupes) before trimming.
404     original_count = len(segments)
405     segments = self._merge_overlapping(segments)
406     if len(segments) < original_count:
407         logger.info(
408             "Merged %d overlapping/duplicate segment(s) before stitching: %d ->
%d.",
409             original_count - len(segments),
410             original_count,
411             len(segments),
412         )
413     if not segments:
414         raise ValueError("No valid (start < end) highlight segments after merge.")
415
416     if not os.path.exists(raw_video_path):
417         raise FileNotFoundError(
418             f"[COMPILER ERROR] Source video file not found: {raw_video_path}"
419         )
420
421     # Defense-in-depth: never let the output name escape output_dir (the API also
422     # validates this at the request boundary).
423     output_path = os.path.join(self.output_dir, os.path.basename(output_filename))
424
425     # Probe video duration so we can detect out-of-range segments
426     video_duration = self._get_video_duration(raw_video_path)
427     if video_duration > 0:
428         logger.info(f"Source video duration: {video_duration:.2f}s")
429     else:
430         logger.warning(
431             "Video duration unknown. Cannot validate segment boundaries against
video length."
432         )
433
434     # We will create temporary clips and stitch them using FFmpeg concat demuxer

```

```

435         temp_clips = []
436         skipped_segments = []
437
438         # Create a temp directory within output_dir
439         with tempfile.TemporaryDirectory(dir=self.output_dir) as tmp_dir:
440             logger.info(
441                 f"Trimming {len(segments)} segments
(begin_padding={self.begin_padding}s, end_padding={self.end_padding}s)..."
442             )
443             # Build the (optional) watermark filtergraph once; it is applied during
every
444             # per-clip re-encode so the mark is baked into each clip identically.
445             watermark_filter = self._watermark_filter(tmp_dir)
446             if watermark_filter:
447                 logger.info("[watermark] applying branding overlay to each clip.")
448             for idx, (raw_start, raw_end) in enumerate(segments):
449                 clip_path, skip_record = self._trim_one_segment(
450                     idx,
451                     raw_start,
452                     raw_end,
453                     segments,
454                     video_duration,
455                     tmp_dir,
456                     raw_video_path,
457                     watermark_filter,
458                 )
459                 if clip_path is not None:
460                     temp_clips.append(clip_path)
461                 if skip_record is not None:
462                     skipped_segments.append(skip_record)
463
464             # Report summary of skipped segments
465             if skipped_segments:
466                 logger.info(
467                     f"\n[SUMMARY] {len(skipped_segments)} out of {len(segments)}
segments were SKIPPED:"
468                 )
469                 for seg_idx, seg_start, seg_end, reason in skipped_segments:
470                     logger.info(
471                         f"    - Segment #{seg_idx + 1} [{seg_start:.2f}s -
{seg_end:.2f}s]: {reason}"
472                     )
473                 logger.info("")
474
475             if not temp_clips:
476                 raise RuntimeError(
477                     f"[COMPILER ERROR] Failed to generate ANY valid highlight clips. "
478                     f"All {len(segments)} segments were skipped. "
479                     f"Check the errors above for details. "
480                     f"Common causes: video file is corrupted, segments are out of range,
or ffmpeg is not installed."
481                 )
482
483             logger.info(
484                 f"Successfully trimmed {len(temp_clips)}/{len(segments)} clips.
Preparing to stitch..."
485             )
486
487             # Create a list file for FFmpeg concat demuxer
488             list_file_path = os.path.join(tmp_dir, "clips.txt")
489             # encoding pinned: a non-ASCII clip path would crash on Windows (cp1252
default);
490             # ffmpeg's concat demuxer reads the list as UTF-8 on every OS.
491             with open(list_file_path, "w", encoding="utf-8") as f:
492                 for clip in temp_clips:

```

```

493             # The concat demuxer resolves relative entries against the list
file's
494             # directory. Use absolute paths so a temp-dir clip does not become
495             # output/tmp.../output/tmp.../clip_0.mp4 on Windows.
496             safe_path = Path(clip).resolve().as_posix().replace("'", "'\\'")
497             f.write(f"file '{safe_path}'\n")
498
499         logger.info("Stitching clips together...")
500         # Stitch them together without re-encoding (since they were encoded with
identical settings)
501         concat_cmd = [
502             ffmpeg_bin(),
503             "-y",
504             "-f",
505             "concat",
506             "-safe",
507             "0",
508             "-i",
509             list_file_path,
510             "-c",
511             "copy",
512             output_path,
513         ]
514
515         try:
516             timeout = ffmpeg_timeout_sec()
517             subprocess.run(
518                 concat_cmd, capture_output=True, check=True, timeout=timeout
519             )
520
521             # Final verification
522             if not os.path.exists(output_path) or os.path.getsize(output_path) == 0:
523                 raise RuntimeError(
524                     f"[COMPILER ERROR] FFmpeg concatenation appeared to succeed, "
525                     f"but output file is missing or empty: {output_path}"
526                 )
527
528             output_size_mb = os.path.getsize(output_path) / (1024 * 1024)
529             logger.info(
530                 f"Dynamic highlight reel saved: {output_path} ({output_size_mb:.1f}
MB)"
531             )
532             return output_path
533         except subprocess.CalledProcessError as e:
534             stderr_msg = e.stderr.decode(errors="replace").strip()
535             raise RuntimeError(
536                 f"[COMPILER ERROR] FFmpeg concatenation failed.\n"
537                 f"  Command: {' '.join(concat_cmd)}\n"
538                 f"  FFmpeg stderr: {stderr_msg[-500:]}"
539             ) from e
540         except subprocess.TimeoutExpired as e:
541             raise RuntimeError(
542                 f"[COMPILER ERROR] FFmpeg concatenation timed out after
{timeout_label(timeout)}.\n"
543                 f"  Command: {' '.join(concat_cmd)}"
544             ) from e
545

```

6. assistant (2026-07-08T05:54:20.528Z)

The compiler.py I read has hardcoded `libx264` ? no `\_select\_encoder`/`encoder` knob that the task attributes to #521/#525. Let me reconcile: check whether origin/master has moved past my local HEAD, and find where `\_select\_encoder` actually lives.

7. user (2026-07-08T05:54:25.262Z)

===HEAD vs origin/master===

bf97700 docs(runlogs): record live SPA CUJ closing the 2 owner-gated validation-latency cells (§9) (#522)
bf97700 docs(runlogs): record live SPA CUJ closing the 2 owner-gated validation-latency cells (§9) (#522)

===commits on origin/master not in HEAD===

===search _select_encoder===

---local working tree---

8. user (2026-07-08T05:54:25.285Z)

File does not exist. Note: your current working directory is C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-indexer.

9. assistant (2026-07-08T05:54:37.986Z)

Key finding: `\_select\_encoder` doesn't exist anywhere in origin/master, and `PHASE\_PLACEMENT.md` doesn't exist at that path. HEAD == origin/master == bf97700. The task's premises (PR #525 merged NVENC/_select_encoder; PHASE_PLACEMENT.md) may not be in master yet. Let me establish ground truth.

10. user (2026-07-08T05:54:41.967Z)

===COMPUTE_DECOUPLED_SERVING dir===

architecture.svg
README.md
runlogs
TESTING_STRATEGY.md

===runlogs dir===

DEPLOY_REPORT_cloud-serving_2026-07-04.md
DEPLOY_REPORT_cloud-serving_a1-cuj_2026-07-04.md
DEPLOY_REPORT_cloud-serving_a1-e2e_2026-07-04.md
DEPLOY_REPORT_cloud-serving_alpha-alliv1_2026-07-06.md
DEPLOY_REPORT_cloud-serving_cp-rebuild_2026-07-05.md
DEPLOY_REPORT_cloud-serving_gpu-resident-nvdec_2026-07-07.md
DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md
README.md
RUNLOG_cloud-serving_exec-status-fastfail_2026-07-07.md
RUNLOG_truly-local_TEMPLATE.md

===PHASE_PLACEMENT anywhere===

===git log for #521 #525===

8866ffe perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine placement (#521)
e99d0c3 feat(stitch): localize the reel watermark to the UI locale (+ OFL Noto fonts) (#304)
2afbadf fix(dryrun): install the watermark font in bootstrap + track the UI-trigger asks (#292)
8736e2d fix(runner): honor stitching watermark/padding + add --detect-from for hi-res reels (#205)
ffdfde0 Merge pull request #103 from avidullu:claude/watermark-bundled-font
112e75d feat(stitcher): bundle Apache-2.0 default watermark font + log fallback reason
10864f3 feat(stitcher): configurable branding watermark on highlight reels (#101)

11. assistant (2026-07-08T05:55:01.825Z)

#521 is merged (8866ffe "run compile/stitch on the L4 worker; config-driven phase?machine placement") but there's **no #525, no PHASE_PLACEMENT.md, no `\_select\_encoder`, no NVENC** anywhere in the tree. The task describes a state ahead of master. Let me verify #521's actual contents and check for any NVENC/encoder work in unmerged branches.

12. user (2026-07-08T05:55:06.990Z)

```
===is 8866ffe(#521) ancestor of HEAD?===
NO - #521 not in HEAD
```

```
===what #521 changed===
```

```
8866ffe perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521)
```

backend/compute/__init__.py	2 +
backend/compute/base.py	34 +++
backend/compute/cloud_run.py	100 ++++++-
backend/config/__init__.py	16 ++
backend/config/models.py	52 ++++
backend/config/placement.py	59 ++++
backend/config/presets.py	10 +-
backend/orchestration.py	322 ++++++++-----
backend/pipeline/stitcher/compiler.py	94 +++++-
backend/worker/__main__.py	139 +++++++
docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md	169 +++++++
docs/COMPUTE_DECOUPLED_SERVING/README.md	1 +
tests/test_api_cloud_dispatch.py	70 +++++
tests/test_compile_dispatch.py	204 +++++++
tests/test_compiler.py	64 +++++
tests/test_compute_backend.py	111 +++++++
tests/test_config_deployment.py	5 +
tests/test_phase_placement.py	138 +++++++
tests/test_worker.py	175 +++++++

19 files changed, 1710 insertions(+), 55 deletions(-)

```
===NVENC/h264_nvenc anywhere in all branches===
```

```
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:backend/compute/base.py
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:backend/compute/cloud_run.py
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:backend/config/models.py
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:backend/config/placement.py
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:backend/config/presets.py
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:backend/orchestration.py
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:backend/pipeline/stitcher/compiler.py
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:backend/worker/__main__.py
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:deploy/cloudrun/Dockerfile
8866ffe2757e8fb87b84dcd4c4158fa1f42f98f:deploy/digitalocean/gpu_setup.sh
(end nvenc search)
```

13. user (2026-07-08T05:55:07.825Z)

```
===all remote branches===
```

```
origin/HEAD -> origin/master
origin/chore/revendor-hosted-compute-picker
origin/docs/runlog-cloud-run-exec-status-fastfail
origin/docs/validation-latency-live-cuj-2026-07-08
origin/feat/521-compile-on-l4
origin/feat/validation-result-cache
origin/fix/cloud-run-exec-status-fastfail
origin/fix/hosted-local-detection-guardrail
origin/master
```

```
===local branches===
```

```
chore/revendor-hosted-compute-picker
+ claude/compile-reliability
+ claude/wasb-config-validation
* codex/issue-523-sd-reel-plan
+ feat/521-compile-on-l4
+ feat/524-l4-primary
  feat/submit-md5-cache-reuse
  fix/cloud-run-exec-status-fastfail
  fix/hosted-local-detection-guardrail
master
perf/validation-shared-decode-pass
```

14. assistant (2026-07-08T05:55:47.755Z)

Now the picture is clear. ****#521 is unmerged**** ? it lives on `origin/feat/521-compile-on-l4` (commit 8866ffe), which is where `PHASE\_PLACEMENT.md`, `\_select\_encoder`, and the NVENC work actually are. There is no #525 in git. The task treats #521 as merged, but master (bf97700) doesn't contain it. Let me establish the branch relationship and read the real base files from the #521 branch.

15. user (2026-07-08T05:55:52.023Z)

```
===is master(bf97700) ancestor of feat/521-compile-on-l4?===
```

```
NO - #521 branch is behind master
```

```
===merge-base===
```

```
699d85bc0abd6e35d3196a4566dcec23f649394e
```

```
===commits on feat/521 not in master===
```

```
8866ffe perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine placement (#521)
```

```
===commits on master not in feat/521===
```

```
bf97700 docs(runlogs): record live SPA CUJ closing the 2 owner-gated validation-latency cells (§9) (#522)
```

16. user (2026-07-08T05:55:54.472Z)

```
===git remote===
```

```
origin https://github.com/Khelsutra/badminton-highlight-indexer.git (fetch)
```

```
origin https://github.com/Khelsutra/badminton-highlight-indexer.git (push)
```

```
===#521 branch compiler.py: _select_encoder + watermark region===
```

```
26: #: Per-encoder availability memo (#521): probing ffmpeg is a subprocess, so a run's first clip pays
27: #: it and the rest reuse the verdict. Keyed by encoder name; cleared only across processes.
31: def _probe_encoder_available(encoder: str) -> bool:
32:     """Can this ffmpeg build actually RUN ``encoder`` (not merely list it)? Runs a tiny 1-frame
34:     real test encode ? not ``-encoders`` grep ? is what distinguishes "usable" from "will fail at
35:     runtime". Memoized. Any error / timeout ? unavailable (caller falls back to libx264)."""
36:     if encoder in _ENCODER_AVAILABLE:
37:         return _ENCODER_AVAILABLE[encoder]
40:     # budget, or 30 s when unbounded) so a wedged encoder can't hang the whole stitch.
54:         "-c:v",
55:         encoder,
66:     _ENCODER_AVAILABLE[encoder] = ok
77:     encoder: str = "libx264",
85:     # Video encoder for the per-clip trim re-encode (#521). "libx264" (CPU) is the default and
86:     # works everywhere; "h264_nvenc"/"hevc_nvenc" use the L4's HARDWARE encoder for the 4K
87:     # re-encode the watermark forces. The requested encoder is PROBED once per run
88:     # (_select_encoder) and falls back to libx264 if this ffmpeg build can't run it.
89:     self.encoder = (encoder or "libx264").strip() or "libx264"
210: def _watermark_filter(self, tmp_dir: str) -> Optional[str]:
211:     """Build the ffmpeg -vf filtergraph that burns the configured watermark (text watermarking is disabled or has nothing to draw. The concat stage stays -c copy
277: def _select_encoder(self) -> Tuple[List[str], str]:
278:     """Resolve the ``-c:v ?`` args used for EVERY per-clip trim re-encode this run, plus the
279:     chosen encoder name. Chosen ONCE (not per clip) so all clips share one codec and the final
280:     concat can stay ``-c copy``. NVENC (hardware) is the L4 win on the 4K re-encode; if the
281:     requested NVENC encoder can't run in this ffmpeg build (CI / non-GPU box), fall back
```

```

to
282:         libx264 so a misconfig never nukes the reel.""
283:         enc = self.encoder
284:         if enc in ("h264_nvenc", "hevc_nvenc") and not _probe_encoder_available(enc):
286:             "[compile] encoder %r is not runnable in this ffmpeg build ? falling back to
"
287:             "libx264 (CPU). On the L4 worker this should be available.",
290:             enc = "libx264"
291:         if enc == "libx264":
293:             return ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"], enc
297:             ["-c:v", enc, "-preset", "p4", "-cq", "23", "-pix_fmt", "yuv420p"],
310:         watermark_filter: Optional[str],
374:         # Default to libx264 (the historical settings) when a caller invokes
_trim_one_segment
375:         # directly without resolving an encoder ? keeps this method's old behaviour intact.
379:         else ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"]
396:         vf_args = ["-vf", watermark_filter] if watermark_filter else []
400:         # whole reel. Clips stay codec-identical either way, so concat -c copy holds.
458:     def compile_highlights(
470:         IDENTICAL encoder settings (codec, preset, CRF, pixel format) so the final
471:         concat can stream-copy (`-c copy`) without re-encoding. Changing the
521:         watermark_filter = self._watermark_filter(tmp_dir)
522:         if watermark_filter:
524:             # Resolve the video encoder ONCE (probe NVENC / fall back) so every clip shares
one codec
525:             # and the final concat stays -c copy (#521).
526:             video_codec_args, chosen_encoder = self._select_encoder()
527:             if chosen_encoder != "libx264":
528:                 logger.info("[compile] encoding clips with %s (hardware).", chosen_encoder)
538:                 watermark_filter,

```

17. assistant (2026-07-08T05:56:34.044Z)

Now the full picture is clear, and it's important:

- `origin/feat/521-compile-on-l4` (commit 8866ffe, ****unmerged****) has diverged from master: it was cut from 699d85b, and master's `bf97700` (#522, which added the §9 CUJ the task cites) landed *after*. So each branch has one commit the other lacks.
- The base this task depends on ? `\_select\_encoder`, the `encoder` knob, `PHASE\_PLACEMENT.md` ? exists ****only**** on the #521 branch, not on master.

Let me read the real base files before deciding how to proceed. I'll dump the #521-branch versions to scratchpad and read them alongside master's runlog §9 and README.

18. user (2026-07-08T05:56:40.510Z)

dumped OK

```

626 C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-
highlight-indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/compiler_521.py
169 C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-
highlight-indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/PHASE_PLACEMENT_521.md
795 total

```

19. user (2026-07-08T05:56:40.652Z)

```

1      # DEPLOY REPORT ? cloud-serving, validation-latency closure (#512 shared-decode + #513
worker skip-revalidation), 2026-07-07
2
3      > Durable **closure record** for the two validation-latency fixes now live on the
`cloud-serving` path: **#512** (`perf(validators)`: decode `scene_cut`/`blur`/`relevance` in ONE
shared pass instead of ~125 random per-sample seeks across 3 captures) and **#513**
(`fix(serving)`: the worker skips its redundant re-validation when the control-plane already
validated + dispatched, via `--skip-validation`). Both shipped in the live image `rally-
worker@sha256:8045f54e?` (built from `origin/master` @ `f5c9dd9`; #512 `20862d7`, #513 `23c1868`

```

are ancestors). Measures the before/after on 3 clips spread across codecs/resolutions, all against the real serving infra. Project `gen-lang-client-0306026826`, region `asia-southeast1`. Honest-limits (incl. the owner-gated live CUJ) in \$6.

```
4
5     ## 1. TL;DR
6
7     | Stage | Result |
8     |---|---|
9     | Live image carries #512+#513 | ? control-plane rev `00018-qvz` + worker Job both on
`sha256:8045f54e?` (= `f5c9dd9`, ancestors `20862d7`/`23c1868` verified) |
10    | Baseline captured (pre-fix) | ? from prod logs on rev `00017-qbd` (image `09d2edf7` =
`66756ad`, pre-#512/#513) ? no re-run needed |
11    | **#513 worker skip-revalidation** | ? real dispatch-style exec: Video 1 validation
**SKIPPED (0 s)**, runs to **SUCCESS / 13 rallies** ? the blur-reject that failed the live CUJ
is **gone** |
12    | **#512 shared-decode speedup (real L4)** | ? Video 1 re-validation **169.3 s ? 69.7 s
(2.43x)** on identical worker hardware, decision byte-identical (sharpness 11.9 < 20.0 both) |
13    | **#512 speedup (off-box A/B, breadth)** | ? 1080p H.264 **1.95x**, 4K HEVC **3.10x**,
5.3K 10-bit HEVC **3.23x** ; every validator decision unchanged OLD?NEW |
14    | Live control-plane "after" number | ? **live-corroborated 2026-07-08** ($9): rev
`00019` validated a fresh 1.83 GiB / 174 s clip in **183 s** (in the 2?3 min band) + reel
compiled end-to-end. Video-1's exact 412 s?~170 s A/B stays a projection (the live CUJ used a
different, heavier clip). |
15
16    **Bottom line:** the two fixes do exactly what they claim on live serving hardware.
**#513** turns the soft 5.3K clip from a **hard FAIL** (worker re-validated at its default
`min_blur=20`, rejected sharpness 11.9, `rc2`, no reel) into a **PASS with 13 rallies** by
trusting the control-plane's gate. **#512** cuts the frame-sampling validators ~**2?3x** (2.43x
measured on the L4; 1.95?3.23x off-box across codecs) with **zero decision drift** OLD?NEW. The
only number still owner-pending is the live control-plane validation wall on rev 00018 ? its
components are measured and it projects to ~2?3 min from the ~6m52s baseline.
17
18    ## 2. Resources
19    - **Control-plane:** `rally-control-plane` Cloud Run **service**, rev **`rally-control-
plane-00018-qvz`** (100% traffic), image `asia-southeast1-docker.pkg.dev/gen-lang-
client-0306026826/rally/rally-worker@sha256:8045f54ea9fcbf?` (built from `f5c9dd9`),
`DEPLOYMENT_PROFILE=cloud-serving`, CPU-only **2 vCPU**, edge-auth ON (`edge_auth_enabled=true`,
`cf_team_domain=khelsutra.cloudflareaccess.com`).
20    - **Worker:** `rally-worker` Cloud Run **Job**, same image `sha256:8045f54e?`, 8 vCPU /
32 GiB / 1x NVIDIA **L4**, `WASB_WEIGHTS=/gcs/models/wasb_badminton_best.pth.tar`,
`WASB_DECODE_BACKEND=nvdec_resident`, GCSFuse mount of `gs://khelsutra-rally-serving`.
21    - **Baseline (pre-fix) surface:** control-plane rev **`00017-qbd`** / worker image
`sha256:09d2edf7?` (= `66756ad` = #507+#508+#509, both **pre-#512/#513**) ? the source of the
prod baseline logs below (2026-07-07, before the 16:15 rev-00018 cutover).
22    - **Corpus:** Video 1 `gs://khelsutra-rally-
serving/videos/uploads/702636eed76145e88aa59c3000c7e415/Video 1 - Badminton.MP4` (1.41 GiB,
5312x2988 yuv420p10le HEVC, 3019 frames, MD5 `8fe6b719603dd7a7b2b6ef543139a4a3`);
`gs://khelsutra-rally-corpus/videos/golden_1080p/Boxhill_Doubles.mp4` (1080p H.264);
`gs://khelsutra-rally-corpus/videos/golden/GX010128.mp4` (2.18 GiB, 4K HEVC, 45298 frames).
23
24    ## 3. Exact commands
25    ```bash
26    # --- Provenance: prove the live image carries both fixes ---
27    gcloud run services describe rally-control-plane --region asia-southeast1 \
28      --format="value(status.latestReadyRevisionName,
spec.template.spec.containers[0].image)" # 00018-qvz ?@sha256:8045f54e?
29    git merge-base --is-ancestor 20862d7 f5c9dd9 && echo "#512 in live image" # yes
30    git merge-base --is-ancestor 23c1868 f5c9dd9 && echo "#513 in live image" # yes
31
32    # --- #513 proof: worker A/B on the LIVE image (Video 1 already staged in the serving
bucket) ---
33    V1="gs://khelsutra-rally-serving/videos/uploads/702636eed76145e88aa59c3000c7e415/Video 1
- Badminton.MP4"
34    # Run A ? WITH --skip-validation (exactly what orchestration._maybe_dispatch_remote
emits on dispatch):
```



```

35     gcloud run jobs execute rally-worker --region asia-southeast1 --async \
36         --args="^@@^infer@@--video-uri=$V1@@--out-uri=gs://khehsutra-rally-
serving/_valclose/v1-skipval@@--segmenter=wasb_hybrid@@--skip-validation@@--
set=deployment.compute_target=local_gpu@@--set=indexing.detector_impl=native@@--
set=indexing.skip_ai_handoff=true@@--set=indexing.wasb.batch_size=64@@--
set=indexing.wasb.prefetch_batches=4@@--set=indexing.wasb.timeout_sec=0"
37     # Run B ? WITHOUT --skip-validation (worker re-validates with the NEW #512 validators ?
still blur-rejects):
38     gcloud run jobs execute rally-worker --region asia-southeast1 --async \
39         --args="^@@^infer@@--video-uri=$V1@@--out-uri=gs://khehsutra-rally-serving/_valclose/v
1-reval@@--segmenter=wasb_hybrid@@--set=deployment.compute_target=local_gpu@@--
set=indexing.detector_impl=native@@--set=indexing.skip_ai_handoff=true"
40
41     # --- #512 timing: bracket "pulled ? -> in.mp4" ? "validation FAILED/SKIPPED" ?
"Phase 2 (WASB) completed" ---
42     gcloud logging read 'resource.type="cloud_run_job" AND
labels."run.googleapis.com/execution_name"="rally-worker-<exec>" \
43         --freshness 1d --format="value(timestamp, textPayload)"
44
45     # --- #512 off-box A/B (isolated clones; run_all decodes the SAME sample indices in both
trees) ---
46     git clone --local --no-hardlinks <repo> repo_old && git -C repo_old checkout 23c1868 #
pre-#512 (per-sample seeks)
47     git -C repo_old worktree add --detach repo_new f5c9dd9
# live image src (shared pass)
48     python valbench.py <repo_old|repo_new> <video> <label> 2 # times
registry.run_all(video, load_config())
49     ``
50
51     ## 4. Run results
52
53     ### 4a. Video 1 (5312x2988 10-bit HEVC, 3019 frames) ? the clip that motivated #512/#513
54
55     | Metric | BEFORE ? pre-fix (rev `00017`, img `09d2edf7`) | AFTER ? #512+#513 live (rev
`00018`, img `8045f54e`) |
56     |---|---|---|
57     | Control-plane validation | **411?412 s** (6m52s) ? passed (baked `min_blur=8.0`) |
**projected ~170 s (2?3 min)**; **live-corroborated 2026-07-08 ? 183 s** on rev `00019` ($9 ? a
fresh 1.83 GiB clip, not a Video-1 re-run) |
58     | Worker (re-)validation | **169.3 s** ? ? **BLUR-REJECT** (sharpness 11.9 < default
20.0) | **0 s ? SKIPPED** (`--skip-validation`, #513) |
59     | &nbsp;&nbsp;&nbsp;? counterfactual: worker re-val if NOT skipped | *(169.3 s, OLD
validators)* | **69.7 s** (NEW #512 validators, **2.43***, same 11.9<20 decision) |
60     | End-to-end | ~10m31s wall ? **FAILED** (0 rallies, no reel) | worker ? **SUCCESS** in
~2m38s exec ? **13 rallies**; **reel live-confirmed 2026-07-08** ($9 ? fresh-clip CUJ compiled
an 11-clip reel end-to-end) |
61     | Rallies produced | **0** (rejected before detection) | **13** (`report.json`
`status=success`, 13 segments) |
62     | Pass / fail | ? **FAIL** | ? **PASS** |
63
64     **Prod baseline evidence (rev 00017, 2026-07-07, unmodified):** control-plane `13:49:18
Running validations ? 13:56:10 Video passed checks` = 411.4 s (and `14:03:56 ? 14:10:48` = 411.9
s). Worker exec `rally-worker-2lbz2` (dispatched by the 14:03 run, OLD image): `pulled 14:11:35
? validation FAILED: blur_check ? sharpness 11.9 below 20.0 @14:14:24 ? exit(2)` ? a
**successful GPU dispatch that failed after the fact**, the exact CUJ #513 fixes.
65
66     **After evidence (live image `8045f54e`, 2026-07-07):**
67     - `rally-worker-k29rk` (**Run A, --skip-validation**): `pulled 17:09:24 ? [worker]
validation SKIPPED (--skip-validation) 17:09:27 ? native WASB (torch 2.6.0+cu124, cuda True) ?
Phase 2 (WASB) completed in 100.9 s, 13 candidate windows (1510/3019 visible) ? emitted 13
rallies ? exit(0)`. Exec wall 157.7 s; `gpu_active`?100.9 s (WASB);
`bytes_out`=`report.json`+`intervals.csv`. Output
`gs://?/_valclose/v1-skipval/outputs/8fe6b719?/`, `status=success`, 13 segments.
68     - `rally-worker-t8llx` (**Run B, re-validation on the NEW image**): `pulled 17:09:29 ?
validation FAILED: blur_check ? sharpness 11.9 below 20.0 @17:10:39 ? exit(2)`. Re-val wall

```

****69.7 s**** vs the OLD image's 169.3 s = ****2.43×** on identical L4 hardware**, and the ****decision** is byte-identical** (11.9 < 20.0) ? #512 is pure speed, zero behavior change. (This run also shows ***why*** #513 is needed: absent the skip, the worker's default ``min_blur=20`` still rejects the clip the control-plane passed at 8.0.)

69

70 **### 4b. Off-box #512 A/B ? codec/resolution breadth (16-core local box, cv2 4.13.0, warm, `run\_all`)**

71

72 Isolated clones at ``23c1868`` (OLD, per-sample seeks) vs ``f5c9dd9`` (NEW, shared pass); identical sample indices ? a fair A/B. ****Every validator decision is identical OLD?NEW**** on every clip (no PASS/FAIL drift from the fix).

73

	Clip	codec / res	OLD (warm)	NEW (warm)	Speedup	decisions OLD?NEW
74		---	---	---	---	---
75		Boxhill_Doubles	1080p H.264	10.91 s	5.60 s	**1.95×
76		PASS)				6/6 identical (all
77		GX010128	4K HEVC (45298 fr)	35.76 s	11.54 s	**3.10×
78		PASS)				6/6 identical (all
79		Video 1	5.3K 10-bit HEVC	173.20 s	53.62 s	**3.23×
80		PASS?)				6/6 identical (all

80 ? On this Windows/cv2 box the 10-bit-HEVC decode yields a Laplacian variance ****above**** 20 so blur PASSES off-box, whereas the Linux serving stack measures ****11.9**** and rejects. That platform decode-stack difference is ****orthogonal to #512**** (which changes only ***which frames are sampled***, identically in both trees); the authoritative blur verdict and the real speedup come from the L4/prod (§4a). Off-box contributes only the OLD?NEW ****ratio****, which is decode-stack-independent.

81

82 The win grows with per-frame decode cost (random-seek re-decode is what #512 removes), and it ****triangulates**** with the real-L4 Video 1 A/B (2.43×): the off-box ratios bound the control-plane projection.

83

84 **### 4c. Control-plane projection (2-vCPU)**

85 Baseline ****~412 s****. Applying the L4-measured ****2.43×** ? ****~170 s (~2m50s)****; the off-box ratios (1.95?3.23×) bracket ****~2?3 min****. Matches the pre-fix design estimate (~6.5 min ? ~2 min). Exact live number is owner-gated (§6).

86

87 **## 5. Rollout state**

88 - ****Already deployed ? this is a measurement/closure record, not a deploy.**** #512+#513 went live with the ``8045f54e`` image cutover (control-plane rev ``00018-qvz``, worker Job same digest) on 2026-07-07. No infra change was made for this report.

89 - ****Rollback (pre-existing levers):**** control-plane ``gcloud run services update-traffic rally-control-plane --region asia-southeast1 --to-revisions rally-control-plane-00017-qbd=100``; worker ``gcloud run jobs update rally-worker --region asia-southeast1 --image ?/rally-worker@sha256:09d2edf7?``. Both revert to the pre-#512/#513 image.

90 - ****Scratch outputs**** written under ``gs://khelsutra-rally-serving/_valclose/`` (isolated prefix; safe to delete).

91

92 **## 6. Honest limits / not-verified**

93 - ****Live end-to-end CUJ is owner-gated.**** ? ****CLOSED 2026-07-08 ? see §9**** (owner drove one live SPA upload on rev ``00019``; both ? cells now measured). Original limit, retained for context: ``POST /api/process`` on rev ``00018`` requires a Cloudflare-Access JWT (``edge_auth_enabled=true``; a direct call returns ****401**** ``missing Cf-Access-Jwt-Assertion header``). So the two ? cells ? the ****live control-plane validation wall on rev 00018**** and the ****stitched reel**** ? need the owner to drive one upload (or hand over a CF-Access JWT / service token). Everything else here is measured on the real serving infra. ****Recipe to close them:**** upload "Video 1 - Badminton.MP4" via the SPA (compute=cloud), then ``gcloud logging read '?service_name="rally-control-plane"?("Running validations" OR "passed checks")'`` for the wall, and confirm the dispatched worker exec logs ``validation SKIPPED`` + emits rallies + the reel compiles. Expect CP validation ~2?3 min and a PASS.

94 - ****#513's worker win is shown via a dispatch-faithful direct exec****, not a control-plane-initiated dispatch (same auth gate). The ``--skip-validation`` flag, its emission in ``orchestration._maybe_dispatch_remote``, and its handling in ``worker.cmd_infer`` are unit-tested on master; Run A exercises the exact worker code path with the exact flag. ? ****UPDATE 2026-07-08 (§9): now also proven on a genuine control-plane-initiated dispatch**** (``rally-worker-64bhj``):

```

`_maybe_dispatch_remote` emitted `--skip-validation` on a live SPA run and the worker honored it
(validation SKIPPED, 0 re-validation).
95 - **Rallies, not a reel.** Run A used `skip_ai_handoff=true` (secret-free) so the 13
"rallies" are local candidate windows (no Gemini refine, no stitched mp4) ? this validates
decode+detect+windowing end-to-end and that the clip passes; the reel stitch is the owner-
CUJ step above. The 13 segments are its inputs.
96 - Off-box absolute times are from a Windows 16-core box (cv2/FFmpeg), not the 2-vCPU
Linux control-plane ? so the ratios transfer (the change is algorithmic: seeks?shared
forward-hybrid pass), the absolutes don't, and (per ?) the 10-bit blur verdict can differ; the
control-plane absolute is anchored to the prod 412 s baseline, not the local numbers. Video 1
off-box is a cross-check of the authoritative real-L4 2.43x.
97 - #515 (crash-before-marker fast-fail) is NOT in this image (`ef41473` post-dates
`f5c9dd9`); it's unrelated to validation latency and out of scope for this closure.
98
99 ## 7. Cost
100 - 2 L4 worker execs (Run A 157.7 s + Run B 109.6 s ? 268 s of L4 time) ? ~$0.5?0.7.
Off-box A/B: $0 (local). No VM, no image build. GCS read for the 3 downloaded clips (~5.1
GiB) is intra-project.
101
102 ## 8. Closure verdict
103 ? CLOSED ? #512 and #513 are proven on live cloud-serving hardware. #513 converts
the motivating soft-focus 5.3K clip from a post-dispatch FAIL into a 13-rally PASS by
eliminating the worker's redundant re-validation (169 s + wrongful blur-reject ? 0 s); #512
speeds the frame-sampling validators 2.43x on the real L4 and 1.95?3.23x off-box across
codecs, with zero decision drift OLD?NEW on every clip. The last owner-gated residual is
now closed live 2026-07-08 ($9): an owner-driven SPA CUJ on rev `00019` validated a fresh
1.83 GiB / 174 s clip in 183 s (in the projected 2?3 min band) and compiled an 11-clip reel
end-to-end ? with the worker `validation SKIPPED` on a genuine control-plane dispatch. (The
Video-1-specific 412 s?~170 s A/B stays a projection by design ? the live CUJ used a different,
heavier clip.) Validation-latency work is done.
104
105 ## 9. Update (2026-07-08) ? live SPA CUJ closes the two owner-gated cells (on rev
`00019`)  

106
107 The two ? cells in $4a (the live control-plane validation wall and the end-to-end
reel) were owner-gated on the Cloudflare-Access `POST /api/process`. On 2026-07-08 the
owner drove one real SPA upload (compute=Cloud) end-to-end, so both are now measured on live
serving infra. Mirrors the gpu-resident report's "S5 Update" append ? the body above is
unchanged; this section records what the live run added.
108
109 Two honest deltas from the recipe (state moved after this report's 2026-07-07 image
cutover):
110 1. Live rev rolled `00018` ? `00019`. Traffic is 100% on live-control-
plane-00019-96f, image live-worker@sha256:c6393d27b4ae? (tag :latest), deployed
2026-07-07T17:06Z ? ~51 min after the 00018`/8045f54e` cutover $2 documents; the worker Job is
on the same :latest`=c6393d27` digest (both surfaces match). It's an undocumented forward-
bump, but #512 (20862d7`) and #513 (23c1868`) are ancestors of origin/master, f5c9dd9` (the
00018` src) is an ancestor of master, and the run empirically shows both fixes live (one-
pass validation wall + worker validation SKIPPED). Measured on what is actually serving.
111 2. A fresh clip, not Video 1. The upload was "First CUJ.MP4" ? 1.83 GiB
(1,960,043,219 B), source duration 174.17 s (?84 Mbps, 4K-class ? heavier than Video 1's
1.41 GiB), video_id 2f31b45e?. So the live wall corroborates the Video-1 projection on a
comparably-heavy clip; it is not a re-measurement of Video 1's 412 s?~170 s A/B (a different
clip has a different absolute wall).
112
113 Full phase breakdown (UTC 2026-07-08, rev `00019`):
114
115 | # | Phase | Window | Wall | Machine | Notes |
116 |---|---|---|---|---|---|
117 | 1 | Upload (client?server, 80-part multipart) | 04:25:56 ? 04:28:51 | ~2m55s |
client uplink | 1.83 GiB @ ~84 Mbps ? owner bandwidth, not infra |
118 | 2 | Accept ? validation start (stage) | 04:28:52 ? 04:29:22 | ~30 s | control-plane |
stage upload to serving bucket |
119 | 3 | Validation (#512 shared-decode) ? | 04:29:22 ? 04:32:25 | 183 s (3m03s) |
control-plane 2 vCPU | one-pass scene_cut/blur/relevance; PASSED |

```

```

120      | 4 | Dispatch ? worker ready (cold start) | 04:32:25 ? 04:33:20 | ~55 s | Cloud Run Job
| image pull + GCSFuse mount + torch/cuda init |
121      | 5 | **Worker skip-validation (#513)** ? | 04:33:20 ? 04:33:24 | ~4 s | L4 |
`validation SKIPPED`, 0 re-validation |
122      | 6 | Detection / segment (WASB) | 04:33:24 ? 04:36:20 | **161.6 s** | L4 (cuda) | 15
candidate windows, 2561 visible pts |
123      | 7 | Ingest + finalize | 04:36:20 ? 04:36:24 | ~4 s | control-plane | intervals?DB,
`status=success` |
124      | ? | (user reviews rallies, clicks Compile) | 04:36:24 ? 04:37:26 | ~62 s idle | ? |
human-in-the-loop |
125      | 8 | **Compile / stitch** ? slow | 04:37:26 ? 04:50:21 | **~12m55s** | control-plane 2
vCPU | 11 clips, ~67 s/clip 4K re-encode+watermark; concat ~2 s; mirror ~11 s |
126
127      Server pipeline (process?success, excl. upload+idle): **7m32s**. Full CUJ (upload?reel):
**~24m25s** wall.
128
129      **Cell 1 ? control-plane validation AFTER (live):** **183 s (3m03s)** on rev `00019`,
PASSED ? the #512 shared-decode validators on the real 2-vCPU serving path, in the projected
**2?3 min** band on a clip *heavier* than Video 1 (pre-#512 per-sample seeks were the
~6m52s-class baseline). Evidence: `Running validations 04:29:22.789 ? Video passed checks
04:32:25.780`.
130
131      **Cell 2 ? end-to-end reel (live):** `/api/compile` stitched **11 of the 15** detected
rallies (0.5 s pad, per-clip branding watermark) from the 174.17 s source ? **`gs://khehsutra-
rally-serving/highlights/2f31b45e?/first-cuj-ks-highlight-reel.mp4`** (**581.84 MiB**, saved
04:50:10 "Dynamic highlight reel saved (581.8 MB)", mirrored 04:50:21, served via
`/api/highlights` 307). **Reel ? produced end-to-end on the live control-plane** ? source fetch
? per-clip trim+watermark ? concat ? mirror-upload, all exercised.
132
133      **Bonus ? #513 proof upgraded** from the $6 dispatch-faithful *direct exec* to a genuine
**control-plane-initiated dispatch** (`rally-worker-64bhj`): `_maybe_dispatch_remote` emitted
`--skip-validation` on a live SPA run and the worker honored it (0 re-validation).
134
135      **? Finding ? stitch is the scaling bottleneck (follow-ups filed).** The compile ran
**~13 min on the CPU-only 2-vCPU control-plane** (~45?90 s per ~7 s clip: the watermark forces a
full 4K re-encode, no NVENC). With validation (183 s) *also* on the control-plane and serialized
by `_LOCAL_COMPUTE_LANE`, one run pins the control-plane for ~16 min of heavy CPU ? it won't
scale past a few concurrent users. Filed: **#521** (move compile/stitch to the L4 worker + make
phase?machine placement config-driven) and **#520** (step-level debug tracing, server+UI, to a
low-latency sink ? this breakdown had to be hand-correlated across 3 log streams).
136
137      **Cost (heuristic):** one L4 detect exec (`rally-worker-64bhj`, ~222 s) ? **$0.15?0.55**
(rate-dependent, dominant) + control-plane CPU for validation+stitch (~960 s @ 2 vCPU) ?
**~$0.05** + reel egress (581 MB, if downloaded to internet) ? **~$0.07** + storage/ops **<
$0.02** + Gemini **$0** (skip_ai_handoff). **? $0.3?0.7 total**, ~one L4 exec as projected ? but
the stitch's real cost is **concurrency** (~16 min of control-plane heavy CPU), not dollars.
138
139      **Net:** both owner-gated cells are now live-measured on rev `00019` ? validation wall
**183 s** (corroborating the 2?3 min projection) and the reel compiles end-to-end. #512 one-pass
validation and #513 skip-validation-on-real-dispatch are now both directly observed in prod. The
Video-1 412 s?~170 s A/B remains a projection by design (different, heavier clip).
140

```

20. assistant (2026-07-08T05:56:54.058Z)

Now I have the CUJ data: \$9 row 8 shows **compile/stitch ~12m55s** (11 clips × ~67s/clip 4K re-encode forced by the watermark; concat only ~2s). That's the cost this task targets. Let me read the #521-branch compiler (the real base) and PHASE_PLACEMENT.md \$3.

21. user (2026-07-08T05:56:56.970Z)

```

1      import logging
2      import os
3      import subprocess
4      import tempfile

```

```

5     from pathlib import Path
6     from typing import Any, Dict, List, Optional, Tuple, Union
7
8     from backend.utils.ffmpeg import ( # P2a: single ffmpeg/ffprobe resolver
9         ffmpeg_bin,
10        ffmpeg_timeout_sec,
11        ffprobe_bin,
12        ffprobe_timeout_sec,
13        timeout_label,
14    )
15
16    logger = logging.getLogger(__name__)
17
18    # Bundled, commercially-clean default watermark font (Apache-2.0). Shipped so the
overlay
19    # renders out-of-the-box without a system fontconfig setup (a common gap on Windows,
where
20    # `font='Sans'` fails). Overridable via stitching.watermark.font_path. This module lives
at
21    # backend/pipeline/stitcher/, so parents[2] == backend/.
22    _BUNDLED_FONT = (
23        Path(__file__).resolve().parents[2] / "assets" / "fonts" / "RobotoSlab-Regular.ttf"
24    )
25
26    #: Per-encoder availability memo (#521): probing ffmpeg is a subprocess, so a run's
first clip pays
27    #: it and the rest reuse the verdict. Keyed by encoder name; cleared only across
processes.
28    _ENCODER_AVAILABLE: Dict[str, bool] = {}
29
30
31    def _probe_encoder_available(encoder: str) -> bool:
32        """Can this ffmpeg build actually RUN ``encoder`` (not merely list it)? Runs a tiny
1-frame
33        lavfi test encode to null. NVENC is often *compiled in* on boxes without an NVIDIA
GPU, so a
34        real test encode ? not ``-encoders`` grep ? is what distinguishes "usable" from
"will fail at
35        runtime". Memoized. Any error / timeout ? unavailable (caller falls back to
libx264)."""
36        if encoder in _ENCODER_AVAILABLE:
37            return _ENCODER_AVAILABLE[encoder]
38        ok = False
39        # The probe is a tiny 1-frame encode ? bound it to a short timeout (? the configured
ffmpeg
40        # budget, or 30 s when unbounded) so a wedged encoder can't hang the whole stitch.
41        configured = ffmpeg_timeout_sec()
42        probe_timeout = min(30.0, float(configured)) if configured else 30.0
43        try:
44            proc = subprocess.run(
45                [
46                    ffmpeg_bin(),
47                    "-hide_banner",
48                    "-f",
49                    "lavfi",
50                    "-i",
51                    "color=c=black:s=64x64:d=0.1",
52                    "-frames:v",
53                    "1",
54                    "-c:v",
55                    encoder,
56                    "-f",
57                    "null",
58                    "-",
59                ],

```

```

60         capture_output=True,
61         timeout=probe_timeout,
62     )
63     ok = proc.returncode == 0
64 except Exception: # noqa: BLE001 ? a probe failure/timeout just means "not usable ?
fall back"
65     ok = False
66     _ENCODER_AVAILABLE[encoder] = ok
67     return ok
68
69
70 class VideoCompiler:
71     def __init__(
72         self,
73         output_dir: str,
74         end_padding: float = 1.0,
75         begin_padding: float = 0.5,
76         watermark: Optional[Any] = None,
77         encoder: str = "libx264",
78     ):
79         self.output_dir = output_dir
80         self.end_padding = end_padding
81         self.begin_padding = begin_padding
82         # Optional branding overlay. Accepts a WatermarkConfig model or a plain dict;
83         # normalized to a dict (or None) so this module stays config-package-agnostic.
84         self.watermark = self._normalize_watermark(watermark)
85         # Video encoder for the per-clip trim re-encode (#521). "libx264" (CPU) is the
default and
86         # works everywhere; "h264_nvenc"/"hevc_nvenc" use the L4's HARDWARE encoder for
the 4K
87         # re-encode the watermark forces. The requested encoder is PROBED once per run
88         # (_select_encoder) and falls back to libx264 if this ffmpeg build can't run it.
89         self.encoder = (encoder or "libx264").strip() or "libx264"
90         os.makedirs(output_dir, exist_ok=True)
91
92     @staticmethod
93     def _normalize_watermark(watermark: Optional[Any]) -> Optional[Dict[str, Any]]:
94         if watermark is None:
95             return None
96         if hasattr(watermark, "model_dump"):
97             return watermark.model_dump()
98         if isinstance(watermark, dict):
99             return dict(watermark)
100         return None
101
102     @staticmethod
103     def _ff_quote(path: str) -> str:
104         """Make a filesystem path safe to embed inside single quotes in an ffmpeg
105         filtergraph operand (drawtext fontfile/textfile, movie= source): forward
106         slashes (Windows-friendly), escaped single quotes, AND escaped colons.
107
108         The colon MUST be escaped even inside single quotes: ffmpeg's filtergraph
109         option parser treats ':' as the option separator, so a Windows drive-letter
110         path like ``C:/dir/font.ttf`` otherwise truncates the operand at ``C`` and
111         the encode fails with "Invalid argument". This silently disabled the
112         on-by-default bundled-font watermark (and absolute logo paths) on Windows."""
113         return path.replace("\\", "/").replace("'", r"\'").replace(":", r"\:")
114
115     @staticmethod
116     def _parse_time(val: Union[int, float, str]) -> float:
117         """Normalizes a timestamp value to float seconds.
118         Handles: float/int (pass-through), plain numeric strings ("12.4"),
119         and colon-separated strings ("MM:SS", "HH:MM:SS", "H:MM:SS.ms").
120         This mirrors the parse_time logic used in gemini.py to ensure
121         the stitcher can always consume whatever format the indexer produced."""

```

```

122         if isinstance(val, (int, float)):
123             return float(val)
124         if isinstance(val, str):
125             val = val.strip()
126             if ":" in val:
127                 # Handle MM:SS, HH:MM:SS, etc.
128                 parts = val.split(":")
129                 parts.reverse()
130                 total = 0.0
131                 for i, p in enumerate(parts):
132                     try:
133                         total += float(p) * (60**i)
134                     except ValueError:
135                         pass
136                 return total
137             try:
138                 return float(val)
139             except ValueError:
140                 logger.warning(
141                     f"Could not parse timestamp value '{val}' as a number. Defaulting to
0.0."
142                 )
143                 return 0.0
144             logger.warning(
145                 f"Unexpected timestamp type {type(val).__name__} for value '{val}'.
Defaulting to 0.0."
146             )
147             return 0.0
148
149     @staticmethod
150     def _merge_overlapping(
151         segments: List[Tuple[float, float]],
152     ) -> List[Tuple[float, float]]:
153         """Normalize, sort, and merge overlapping/duplicate (start, end) ranges.
154
155         A highlight reel must never replay the same footage. The same rally can land in
156         the segment list more than once ? e.g. two detector runs accumulated for one
157         video (add_video is an UPSERT, so re-running without a prune leaves both sets),
158         or
159         chunk-overlap duplicates near boundaries. Stitching those verbatim plays the
160         rally
161         twice. Merging any range that overlaps or abuts the previous one collapses true
162         duplicates/overlaps while preserving genuinely distinct (gapped) rallies.
163         Degenerate rows (end <= start) are dropped. Timestamps may be float/int/str.
164         """
165         parsed = []
166         for raw_s, raw_e in segments:
167             s = VideoCompiler._parse_time(raw_s)
168             e = VideoCompiler._parse_time(raw_e)
169             if e > s:
170                 parsed.append((s, e))
171         parsed.sort()
172         merged: List[Tuple[float, float]] = []
173         for s, e in parsed:
174             if merged and s <= merged[-1][1]: # overlaps/abuts previous -> extend it
175                 merged[-1] = (merged[-1][0], max(merged[-1][1], e))
176             else:
177                 merged.append((s, e))
178         return merged
179
180     def _get_video_duration(self, video_path: str) -> float:
181         """Probes the video file to get its total duration in seconds.
182         Returns 0.0 if it cannot be determined."""
183         try:
184             timeout = ffprobe_timeout_sec()

```

```

183         result = subprocess.run(
184             [
185                 ffmpeg_bin(),
186                 "-v",
187                 "error",
188                 "-show_entries",
189                 "format=duration",
190                 "-of",
191                 "default=noprint_wrappers=1:nokey=1",
192                 video_path,
193             ],
194             capture_output=True,
195             text=True,
196             check=True,
197             timeout=timeout,
198         )
199         return float(result.stdout.strip())
200     except subprocess.TimeoutExpired:
201         logger.warning(
202             "Could not determine video duration via ffmpeg: timed out after %s",
203             timeout_label(ffmpeg_timeout_sec()),
204         )
205         return 0.0
206     except Exception as e:
207         logger.warning(f"Could not determine video duration via ffmpeg: {e}")
208         return 0.0
209
210 def _watermark_filter(self, tmp_dir: str) -> Optional[str]:
211     """Build the ffmpeg -vf filtergraph that burns the configured watermark (text
212     and/or logo PNG) into a clip during the per-clip re-encode. Returns None when
213     watermarking is disabled or has nothing to draw. The concat stage stays -c copy
214     because the mark is already baked into every clip identically."""
215     wm = self.watermark
216     if not wm or not wm.get("enabled"):
217         return None
218
219     text = (wm.get("text") or "").strip()
220     logo_path = (wm.get("logo_path") or "").strip()
221     if logo_path and not os.path.exists(logo_path):
222         logger.warning(
223             f"[watermark] logo not found: {logo_path} ? skipping logo overlay."
224         )
225         logo_path = ""
226     if not text and not logo_path:
227         return None
228
229     margin = int(wm.get("margin", 24))
230     opacity = float(wm.get("opacity", 0.85))
231     position = str(wm.get("position", "bottom-right"))
232     # drawtext uses lowercase w/h (main) + tw/th (text); overlay uses W/H (main) +
233     w/h (logo).
234     dt_xy = {
235         "bottom-right": f"x=w-tw-{margin}:y=h-th-{margin}",
236         "bottom-left": f"x={margin}:y=h-th-{margin}",
237         "top-right": f"x=w-tw-{margin}:y={margin}",
238         "top-left": f"x={margin}:y={margin}",
239     }.get(position, f"x=w-tw-{margin}:y=h-th-{margin}")
240     ov_xy = {
241         "bottom-right": f"W-w-{margin}:H-h-{margin}",
242         "bottom-left": f"{margin}:H-h-{margin}",
243         "top-right": f"W-w-{margin}:{margin}",
244         "top-left": f"{margin}:{margin}",
245     }.get(position, f"W-w-{margin}:H-h-{margin}")
246     drawtext = ""

```



```

247         if text:
248             ratio = float(wm.get("font_size_ratio") or 0.035)
249             divisor = max(1, round(1.0 / ratio))
250             # textfile= sidesteps the notorious drawtext text= escaping for arbitrary
strings.
251             tf_path = os.path.join(tmp_dir, "watermark.txt")
252             with open(tf_path, "w", encoding="utf-8") as fh:
253                 fh.write(text)
254             # Font precedence: explicit font_path ? bundled Apache-2.0 default (renders
255             # without fontconfig) ? fontconfig family (last-resort, environment-
dependent).
256             font_path = (wm.get("font_path") or "").strip()
257             if not font_path and _BUNDLED_FONT.exists():
258                 font_path = str(_BUNDLED_FONT)
259             if font_path:
260                 font_opt = f"fontfile='{self._ff_quote(font_path)}'"
261             else:
262                 font_opt = f"font='{wm.get('font', 'Sans')}'"
263             box = ""
264             if wm.get("box", True):
265                 box = f":box=1:boxcolor=black@{min(opacity, 0.5):.2f}:boxborderw=10"
266             drawtext = (
267                 f"drawtext=textfile='{self._ff_quote(tf_path)}':{font_opt}"
268                 f":fontcolor=white@{opacity:.2f}:fontsize=h/{divisor}:{dt_xy}{box}"
269             )
270
271             if logo_path and drawtext:
272                 return
273             f"movie='{self._ff_quote(logo_path)}'[wm];[in][wm]overlay={ov_xy},{drawtext}"
274             if logo_path:
275                 return f"movie='{self._ff_quote(logo_path)}'[wm];[in][wm]overlay={ov_xy}"
276             return drawtext
277
278         def _select_encoder(self) -> Tuple[List[str], str]:
279             """Resolve the ``-c:v ?`` args used for EVERY per-clip trim re-encode this run,
plus the
280             chosen encoder name. Chosen ONCE (not per clip) so all clips share one codec and
the final
281             concat can stay ``-c copy``. NVENC (hardware) is the L4 win on the 4K re-encode;
if the
282             requested NVENC encoder can't run in this ffmpeg build (CI / non-GPU box), fall
back to
283             libx264 so a misconfig never nukes the reel."""
284             enc = self.encoder
285             if enc in ("h264_nvenc", "hevc_nvenc") and not _probe_encoder_available(enc):
286                 logger.warning(
287                     "[compile] encoder %r is not runnable in this ffmpeg build ? falling
back to "
288                     "libx264 (CPU). On the L4 worker this should be available.",
289                     enc,
290                 )
291             enc = "libx264"
292             if enc == "libx264":
293                 # Unchanged historical CPU settings (byte-compatible with the pre-#521
stitch).
294                 return ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"], enc
295                 # NVENC (H.264/HEVC): hardware encode. -cq is the CRF-equivalent constant-
quality knob; p4 is
296                 # a balanced speed/quality preset. yuv420p keeps the output broadly playable.
297                 return (
298                     ["-c:v", enc, "-preset", "p4", "-cq", "23", "-pix_fmt", "yuv420p"],
299                     enc,
300                 )
301
302         def _trim_one_segment(

```

```

302         self,
303         idx: int,
304         raw_start: Union[int, float, str],
305         raw_end: Union[int, float, str],
306         segments: List[Tuple[float, float]],
307         video_duration: float,
308         tmp_dir: str,
309         raw_video_path: str,
310         watermark_filter: Optional[str],
311         video_codec_args: Optional[List[str]] = None,
312     ) -> Tuple[Optional[str], Optional[Tuple[int, float, float, str]]]:
313         """Trim a single segment into its own clip file (extracted verbatim from the
314         per-segment body of compile_highlights). Returns (clip_path, skip_record): on
315         success the clip path and None; on a skip (invalid/out-of-range/ffmpeg failure)
316         None and the (idx, start, end, reason) record to record in skipped_segments."""
317         # Normalize timestamps to float seconds ? handles float, int,
318         # plain numeric strings, and colon-separated formats like "MM:SS"
319         start = self._parse_time(raw_start)
320         end = self._parse_time(raw_end)
321         segment_label = (
322             f"Segment #{idx + 1}/{len(segments)} [{start:.2f}s - {end:.2f}s]"
323         )
324
325         # Validate the segment times
326         if start < 0:
327             logger.warning(
328                 f"{segment_label}: start_time is negative ({start:.2f}s). Clamping to
329                 0."
330             )
331             start = 0.0
332
333         if start >= end:
334             logger.error(
335                 f"{segment_label}: start_time ({start:.2f}s) >= end_time ({end:.2f}s).
336                 Skipping invalid segment."
337             )
338             return None, (idx, start, end, "start >= end")
339
340         # Check if segment extends beyond video duration
341         if video_duration > 0:
342             if start >= video_duration:
343                 logger.error(
344                     f"{segment_label}: start_time ({start:.2f}s) is beyond the video
345                     duration ({video_duration:.2f}s). "
346                     f"This segment cannot be extracted ? the video ran out. Skipping."
347                 )
348                 return None, (idx, start, end, "start beyond video end")
349
350             if end > video_duration:
351                 original_end = end
352                 end = video_duration
353                 logger.warning(
354                     f"{segment_label}: end_time ({original_end:.2f}s) exceeds video
355                     duration ({video_duration:.2f}s). "
356                     f"Clamping end_time to {video_duration:.2f}s. The segment may be
357                     truncated."
358                 )
359
360         clip_path = os.path.join(tmp_dir, f"clip_{idx}.mp4")
361
362         # Precise sub-second trim using fast re-encoding to preserve stream headers
363         padded_start = max(0.0, start - self.begin_padding)
364         padded_end = end + self.end_padding
365
366         # Clamp padded_end to video duration if known

```

```

362         if video_duration > 0 and padded_end > video_duration:
363             padded_end = video_duration
364             logger.info(
365                 f"{segment_label}: Padded end ({end + self.end_padding:.2f}s) exceeds
video duration. "
366                 f"Clamping to {video_duration:.2f}s (end padding partially applied)."
367             )
368
369         trim_duration = padded_end - padded_start
370         logger.info(
371             f"Trimming {segment_label} -> [{padded_start:.3f}s - {padded_end:.3f}s]
(duration: {trim_duration:.2f}s)"
372         )
373
374         # Default to libx264 (the historical settings) when a caller invokes
_trim_one_segment
375         # directly without resolving an encoder ? keeps this method's old behaviour
intact.
376         codec_args = (
377             video_codec_args
378             if video_codec_args is not None
379             else ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"]
380         )
381         base_cmd = [
382             ffmpeg_bin(),
383             "-y",
384             "-ss",
385             str(round(padded_start, 3)),
386             "-to",
387             str(round(padded_end, 3)),
388             "-i",
389             raw_video_path,
390             *codec_args,
391             "-c:a",
392             "aac",
393             "-b:a",
394             "128k",
395         ]
396         vf_args = ["-vf", watermark_filter] if watermark_filter else []
397
398         # Try with the watermark first; if that encode fails (e.g. a bad font/logo
399         # path), retry once WITHOUT it so a branding misconfig can never nuke the
400         # whole reel. Clips stay codec-identical either way, so concat -c copy holds.
401         attempts = [vf_args, []] if vf_args else [[]]
402         produced = False
403         last_exc: Optional[subprocess.CalledProcessError] = None
404         cmd = base_cmd + [clip_path]
405         timeout = ffmpeg_timeout_sec()
406         for attempt_vf in attempts:
407             cmd = base_cmd + attempt_vf + [clip_path]
408             try:
409                 subprocess.run(
410                     cmd, capture_output=True, check=True, timeout=timeout
411                 )
412             except subprocess.TimeoutExpired:
413                 reason = f"ffmpeg timeout after {timeout_label(timeout)}"
414                 logger.error(f"{segment_label}: FFmpeg trim timed out.")
415                 print(f"  Command: {' '.join(cmd)}")
416                 print(f"  {reason}")
417                 return None, (idx, start, end, reason)
418             except subprocess.CalledProcessError as e:
419                 last_exc = e
420                 if attempt_vf and vf_args:
421                     _err = (
422                         e.stderr.decode(errors="replace").strip()[-300:]

```

```

423             if getattr(e, "stderr", None)
424             else ""
425         )
426         logger.warning(
427             f"{segment_label}: watermark encode failed; retrying without it."
428             f"ffmpeg: {_err or '(no stderr captured)'}"
429         )
430         continue
431     break
432     if os.path.exists(clip_path) and os.path.getsize(clip_path) > 0:
433         produced = True
434         if not attempt_vf and vf_args:
435             logger.warning(
436                 f"{segment_label}: encoded WITHOUT the watermark (the watermark
filter failed).")
437             )
438             break
439             last_exc = None # ffmpeg succeeded but produced nothing
440             break
441
442     if produced:
443         return clip_path, None
444     elif last_exc is not None:
445         stderr_msg = last_exc.stderr.decode(errors="replace").strip()
446         logger.error(f"{segment_label}: FFmpeg trim failed.")
447         print(f"  Command: {' '.join(cmd)}")
448         print(
449             f"    FFmpeg stderr: {stderr_msg[-500:]}")
450         ) # Last 500 chars to avoid flooding
451         return None, (idx, start, end, f"ffmpeg error: {stderr_msg[-200:]}")
452     else:
453         logger.error(
454             f"{segment_label}: FFmpeg exited successfully but output clip is missing
or empty. Skipping."
455         )
456         return None, (idx, start, end, "empty output clip")
457
458     def compile_highlights(
459         self,
460         raw_video_path: str,
461         segments: List[Tuple[float, float]],
462         output_filename: str,
463     ) -> str:
464         """
465         Extracts selected segments from a raw video and stitches them together
466         into a seamless highlights file, applying end_padding to prevent abrupt endings.
467
468         .. warning::
469             This method uses a two-stage codec contract: per-clip re-encodes must use
470             IDENTICAL encoder settings (codec, preset, CRF, pixel format) so the final
471             concat can stream-copy (``-c copy``) without re-encoding. Changing the
472             clip encode settings without updating the concat assumptions will produce
473             broken output.
474         """
475         if not segments:
476             raise ValueError("No highlight segments provided to compile.")
477
478         # Never stitch the same footage twice: collapse overlapping/duplicate ranges
479         # (e.g. accumulated re-run segments or chunk-overlap dupes) before trimming.
480         original_count = len(segments)
481         segments = self._merge_overlapping(segments)
482         if len(segments) < original_count:
483             logger.info(
484                 "Merged %d overlapping/duplicate segment(s) before stitching: %d ->

```

```

%d.",
485         original_count - len(segments),
486         original_count,
487         len(segments),
488     )
489 if not segments:
490     raise ValueError("No valid (start < end) highlight segments after merge.")
491
492 if not os.path.exists(raw_video_path):
493     raise FileNotFoundError(
494         f"[COMPILER ERROR] Source video file not found: {raw_video_path}"
495     )
496
497 # Defense-in-depth: never let the output name escape output_dir (the API also
498 # validates this at the request boundary).
499 output_path = os.path.join(self.output_dir, os.path.basename(output_filename))
500
501 # Probe video duration so we can detect out-of-range segments
502 video_duration = self._get_video_duration(raw_video_path)
503 if video_duration > 0:
504     logger.info(f"Source video duration: {video_duration:.2f}s")
505 else:
506     logger.warning(
507         "Video duration unknown. Cannot validate segment boundaries against
video length."
508     )
509
510 # We will create temporary clips and stitch them using FFmpeg concat demuxer
511 temp_clips = []
512 skipped_segments = []
513
514 # Create a temp directory within output_dir
515 with tempfile.TemporaryDirectory(dir=self.output_dir) as tmp_dir:
516     logger.info(
517         f"Trimming {len(segments)} segments
(begin_padding={self.begin_padding}s, end_padding={self.end_padding}s)..."
518     )
519     # Build the (optional) watermark filtergraph once; it is applied during
every
520     # per-clip re-encode so the mark is baked into each clip identically.
521     watermark_filter = self._watermark_filter(tmp_dir)
522     if watermark_filter:
523         logger.info("[watermark] applying branding overlay to each clip.")
524     # Resolve the video encoder ONCE (probe NVENC / fall back) so every clip
shares one codec
525     # and the final concat stays -c copy (#521).
526     video_codec_args, chosen_encoder = self._select_encoder()
527     if chosen_encoder != "libx264":
528         logger.info("[compile] encoding clips with %s (hardware).",
chosen_encoder)
529     for idx, (raw_start, raw_end) in enumerate(segments):
530         clip_path, skip_record = self._trim_one_segment(
531             idx,
532             raw_start,
533             raw_end,
534             segments,
535             video_duration,
536             tmp_dir,
537             raw_video_path,
538             watermark_filter,
539             video_codec_args,
540         )
541         if clip_path is not None:
542             temp_clips.append(clip_path)
543         if skip_record is not None:

```

```

544         skipped_segments.append(skip_record)
545
546     # Report summary of skipped segments
547     if skipped_segments:
548         logger.info(
549             f"\n[SUMMARY] {len(skipped_segments)} out of {len(segments)}
segments were SKIPPED:"
550         )
551         for seg_idx, seg_start, seg_end, reason in skipped_segments:
552             logger.info(
553                 f" - Segment #{seg_idx + 1} [{seg_start:.2f}s -
{seg_end:.2f}s]: {reason}"
554             )
555             logger.info("")
556
557     if not temp_clips:
558         raise RuntimeError(
559             f"[COMPILER ERROR] Failed to generate ANY valid highlight clips. "
560             f"All {len(segments)} segments were skipped. "
561             f"Check the errors above for details. "
562             f"Common causes: video file is corrupted, segments are out of range,
or ffmpeg is not installed."
563         )
564
565     logger.info(
566         f"Successfully trimmed {len(temp_clips)}/{len(segments)} clips.
Preparing to stitch..."
567     )
568
569     # Create a list file for FFmpeg concat demuxer
570     list_file_path = os.path.join(tmp_dir, "clips.txt")
571     # encoding pinned: a non-ASCII clip path would crash on Windows (cp1252
default);
572     # ffmpeg's concat demuxer reads the list as UTF-8 on every OS.
573     with open(list_file_path, "w", encoding="utf-8") as f:
574         for clip in temp_clips:
575             # The concat demuxer resolves relative entries against the list
file's
576             # directory. Use absolute paths so a temp-dir clip does not become
577             # output/tmp.../output/tmp.../clip_0.mp4 on Windows.
578             safe_path = Path(clip).resolve().as_posix().replace("'", "'\\'")
579             f.write(f"file '{safe_path}'\n")
580
581     logger.info("Stitching clips together...")
582     # Stitch them together without re-encoding (since they were encoded with
identical settings)
583     concat_cmd = [
584         ffmpeg_bin(),
585         "-y",
586         "-f",
587         "concat",
588         "-safe",
589         "0",
590         "-i",
591         list_file_path,
592         "-c",
593         "copy",
594         output_path,
595     ]
596
597     try:
598         timeout = ffmpeg_timeout_sec()
599         subprocess.run(
600             concat_cmd, capture_output=True, check=True, timeout=timeout
601         )

```

```

602
603         # Final verification
604         if not os.path.exists(output_path) or os.path.getsize(output_path) == 0:
605             raise RuntimeError(
606                 f"[COMPILER ERROR] FFmpeg concatenation appeared to succeed, "
607                 f"but output file is missing or empty: {output_path}"
608             )
609
610         output_size_mb = os.path.getsize(output_path) / (1024 * 1024)
611         logger.info(
612             f"Dynamic highlight reel saved: {output_path} ({output_size_mb:.1f}
613         MB)"
614         )
615         return output_path
616     except subprocess.CalledProcessError as e:
617         stderr_msg = e.stderr.decode(errors="replace").strip()
618         raise RuntimeError(
619             f"[COMPILER ERROR] FFmpeg concatenation failed.\n"
620             f"    Command: {' '.join(concat_cmd)}\n"
621             f"    FFmpeg stderr: {stderr_msg[-500:]}"
622         ) from e
623     except subprocess.TimeoutExpired as e:
624         raise RuntimeError(
625             f"[COMPILER ERROR] FFmpeg concatenation timed out after
626             {timeout_label(timeout)}.\n"
627             f"    Command: {' '.join(concat_cmd)}"
628         ) from e

```

22. user (2026-07-08T05:56:57.310Z)

```

1      # Phase ? machine placement + compile-on-L4 (#521)
2
3      > **Status:** ? **SHIPPED (code)** . **Owner:** Avi . **Created:** 2026-07-08 .
**Issue:** [#521](https://github.com/Khelsutra/badminton-highlight-indexer/issues/521)
4      > **Relation:** extends [COMPUTE_DECOUPLED_SERVING](./README.md) (the `ComputeBackend`
dispatch seam) and **refines D16(b)** (NVENC-encode-deferred-behind-a-config-knob) now that the
2026-07-08 live CUJ supplies the cost data D16(b) said to revisit on.
5      > **Evidence:** `runlogs/DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md`
S9 (2026-07-08 live CUJ, rev `rally-control-plane-00019`).
6
7      ---
8
9      ## 0. TL;DR
10
11      The compile/stitch phase used to run **in-process on the CPU-only 2-vCPU control-
plane**, where the
12      2026-07-08 live CUJ measured **~12m55s** for an 11-clip / ~80 s reel from a 1.83 GiB 4K
source
13      (~45?90 s per ~7 s clip ? the per-clip branding watermark forces a full 4K re-encode
with **no
14      NVENC**). Validation (183 s) runs there too; both serialize on `_LOCAL_COMPUTE_LANE`, so
one run pins
15      the control-plane for ~16 min ? the API can't scale past a handful of concurrent users.
16
17      #521 makes **where each pipeline phase runs a config edit, not a code change**, and
moves the stitch
18      to the **L4 worker (NVENC + 8 vCPU)** the way detection already dispatches:
19
20      1. **`deployment.phase_placement`** ? a per-phase knob (`validation` / `segment` /
`compile`) with
21          values `control_plane | cloud_run_l4`. Unset ? **today's placement, byte-identical**.
22      2. **`compile ? L4`** ? a `kind='compile'` Cloud Run dispatch mirroring the detect path:
the
23          control-plane resolves the reel's segments, stages a compile-spec to the bucket,

```

```

dispatches
24     `python -m backend.worker compile` on the L4, bucket-polls a done-marker, and returns
the same
25     `{status, filepath, download_url}` shape the SPA already expects.
26     3. **NVENC on the L4** ? `deployment.cloud_stitch_encoder` (default `h264_nvenc`) is
forwarded onto
27     the worker's `stitching.encoder`; libx264 stays the everywhere-safe fallback (auto-
probed).
28
29     The cloud-serving preset opts compile onto the L4
(`phase_placement.compile=cloud_run_l4`).
30
31     ---
32
33     ## 1. The placement model
34
35     `resolve_phase_placement(cfg, phase)` (in `backend/config/placement.py`) is the **single
source of
36     truth3 ? orchestration and the worker never re-derive the mapping. Precedence: an
explicit
37     `deployment.phase_placement.<phase>` wins; otherwise the **legacy default** reproduces
pre-#521
38     behaviour exactly:
39
40     | phase | knob unset (legacy default) | meaning |
41     |---|---|---|
42     | `validation` | `control_plane` | the authoritative gate (#512/#513) stays on the
control-plane |
43     | `segment` (detect) | `cloud_run_l4` **iff** `compute_target==cloud_run`, else
`control_plane` | today's DETECT dispatch |
44     | `compile` (stitch) | `control_plane` | the in-process ffmpeg stitch |
45
46     So a config with **no** `phase_placement` block behaves exactly as before this existed.
To relocate a
47     phase, set its knob:
48
49     ```jsonc
50     // Move the stitch to the L4 (what the cloud-serving preset does):
51     { "deployment": { "compute_target": "cloud_run",
52                       "phase_placement": { "compile": "cloud_run_l4" } } }
53
54     // Force detection back onto the control-plane on a cloud box (e.g. a debugging run):
55     { "deployment": { "compute_target": "cloud_run",
56                       "phase_placement": { "segment": "control_plane" } } }
57
58     // Move validation off the control-plane too (must pair with a cloud segment ? see §4):
59     { "deployment": { "compute_target": "cloud_run",
60                       "phase_placement": { "validation": "cloud_run_l4" } } }
61     ```
62
63     `control_plane` runs the phase in-process on the (CPU-only) control-plane service;
`cloud_run_l4`
64     dispatches it to the GPU/NVENC Cloud Run worker Job.
65
66     ---
67
68     ## 2. compile ? L4 dispatch (the flow)
69
70     `orchestration._execute_compile` consults `resolve_phase_placement(cfg, "compile")`.
When it resolves
71     to `cloud_run_l4` **and** cloud dispatch is available (`_cloud_available`), it calls
72     `_dispatch_compile_remote`; otherwise it stitches in-process via the shared
`_stitch_reel` core
73     (unchanged).
74

```



```

75     `_dispatch_compile_remote` (mirrors `_maybe_dispatch_remote`):
76
77     1. **Resolve** the reel's segments HERE ? the control-plane owns the DB
78     (`_resolve_compile` ?
79         `(video, kept, out_name)`), applying the `stitching.min_shot_count` floor as before.
80     2. **Stage a compile-spec** to `**no control-plane DB** ? everything rides the spec.
86     3. **Dispatch** a `compile` Cloud Run Job execution (`CompileJobSpec` ?
87         `CloudRunCompute.run_compile`)
88         running `python -m backend.worker compile --spec-uri ? --done-uri ? --set
89         storage.output_root=?`,
90         forwarding the bucket coordinates (the worker Job has no `DEPLOYMENT_PROFILE`).
91     4. **Bucket-poll** the done-marker at `**shared
98         `_stitch_reel`
99         core** (no forked pipeline), mirror-uploads it to
100         `**Fallback:** if compile is placed on the L4 but the box isn't actually cloud-wired (no
109     output bucket
110     / Cloud Run env), `_dispatch_compile_remote` logs and stitches in-process ? a slow reel
111     beats no reel.
112
113     ---
114
115     ## 3. NVENC on the L4 ? and the D5/D16(b) parity note
116
117     The dominant term in the CUJ stitch was the **per-clip 4K re-encode** the branding
118     watermark forces
119     (~67 s/clip on 2 vCPU with libx264). The L4 has a hardware H.264/HEVC encoder, so #521
120     adds an encoder
121     knob:
122
123     - `StitchingConfig.encoder` ? `libx264` (default, CPU, works everywhere incl. CI) |
124     `h264_nvenc` |
125     `hevc_nvenc`. `VideoCompiler` **probes** the requested encoder once per run (a 1-frame
126     lavfi test
127     encode) and **falls back to libx264** if this ffmpeg build can't run it, so a
128     misconfig / non-GPU
129     box never breaks. All clips share the one chosen encoder ? the concat stays `-c copy`.
130     - `deployment.cloud_stitch_encoder` (default `h264_nvenc`) ? what the compile-dispatch
131     forwards onto
132     the worker's `stitching.encoder` for the L4 run. Set it to `libx264` to run the L4
133     stitch on CPU.
134     - The control-plane's **own** local stitch always uses `stitching.encoder` (libx264 by
135     default) ? the
136     CP has no GPU.
137
138     **Parity note (honest):** [README.md](./README.md) **D5/D16(b)** kept stitch **encode =
139     libx264** to

```

```

118     protect a byte-exact stitch parity guarantee, and explicitly **deferred NVENC encode
behind a config
119     knob, "revisit only if M7 cost data justifies relaxing parity."** The 2026-07-08 CUJ
(~13 min stitch,
120     a hard scaling wall past ~5 concurrent users) **is** that cost data. #521 lands exactly
the deferred
121     shape ? NVENC **behind a config knob**, libx264 the default, config-reversible per
phase. Enabling
122     NVENC relaxes byte-identity **at the reel ENCODE layer only** (a cosmetic, shareable
artifact); the
123     **WINDOW/rally-level parity** (the analytical output) is untouched ? detection and
windowing are
124     unchanged. If bit-exact reels are ever required, set `cloud_stitch_encoder=libx264`.
125
126     > Follow-up (tracked, not in #521's first cut): a **pre-rendered / overlay-filter
watermark** so the
127     > stitch isn't a full re-encode at all (the deeper win beyond swapping the encoder).
128
129     ---
130
131     ## 4. Validation placement + the "nobody validates" guard
132
133     `phase_placement.validation` moves the CPU-bound validator suite off the 2-vCPU control-
plane too,
134     via the existing `--skip-validation` lever (#513):
135
136     - **`control_plane`** (default): the control-plane runs the validators and is the
authoritative gate;
137     the worker **skips** its redundant re-validation (`skip_validation=True`) ? today's
#512/#513
138     behaviour, byte-identical.
139     - **`cloud_run_l4`**: when detection **actually** dispatches to the L4, the control-
plane **delegates**
140     ? it does NOT run the validator suite (`execute_processing` skips it), and the
dispatch carries
141     **no** `--skip-validation`, so the **worker** validates and rejects a bad clip via a
non-zero rc
142     (surfaced as a failed job).
143
144     **Footgun guard:** a clip can only be validated on the worker if the worker also
**detects** it there.
145     Two protections:
146
147     1. **Config-load guard** (`DeploymentConfig` model-validator): `validation=cloud_run_l4`
requires the
148     effective segment placement to also be `cloud_run_l4` (explicit, or via
`compute_target=cloud_run`)
149     ? else config load fails with an actionable error.
150     2. **Runtime guard** (`_will_detect_on_cloud`): if a per-request `compute='local'`
forces detection
151     in-process (no worker), the control-plane **still validates** ? never leave a clip
unvalidated
152     because nobody ran the gate.
153
154     ---
155
156     ## 5. Files & tests
157
158     | Area | Code | Tests |
159     |---|---|---|
160     | Placement resolver + config | `backend/config/placement.py`,
`backend/config/models.py` (`PhasePlacementConfig`, `DeploymentConfig.phase_placement` /
`cloud_stitch_encoder`, guard) | `tests/test_phase_placement.py`,
`tests/test_config_deployment.py` |
161     | Preset | `backend/config/presets.py` (`cloud-serving` ?

```

```

`phase_placement.compile=cloud_run_l4`) | `tests/test_phase_placement.py` |
162 | compile dispatch | `backend/orchestration.py` (`_execute_compile`,
`_dispatch_compile_remote`, `_stitch_reel`, `_will_detect_on_cloud`) |
`tests/test_compile_dispatch.py` |
163 | compute backend | `backend/compute/base.py` (`CompileJobSpec`, `run_compile`),
`backend/compute/cloud_run.py` (`run_compile`, `_await_marker`) |
`tests/test_compute_backend.py` |
164 | worker | `backend/worker/__main__.py` (`compile` subcommand, `cmd_compile`) |
`tests/test_worker.py` |
165 | NVENC encoder | `backend/pipeline/stitcher/compiler.py` (`encoder`, `_select_encoder`,
`_probe_encoder_available`) | `tests/test_compiler.py` |
166 | validation delegation | `backend/orchestration.py` (`maybe_dispatch_remote` skip
flag, `_execute_processing`) | `tests/test_api_cloud_dispatch.py` |
167
168 All GPU-free + cloud-free: a fake `CloudRunJobClient` simulates the worker (reads the
staged spec,
169 writes the reel + done-marker into a local "bucket"), mirroring
`tests/test_api_cloud_dispatch.py`.
170

```

23. assistant (2026-07-08T05:57:56.011Z)

I now have full context. Before writing any code, I need to surface a material discrepancy between the task's premise and the actual repo state, because it determines the base branch ? and that's a fork I shouldn't pick silently given the repo's "branch off master / NO PR stacking" rule.

****What I found:****

- ****#521 is *not* merged to master.**** It lives on `origin/feat/521-compile-on-l4` (commit 8866ffe), which has *diverged* from master: it was cut from `699d85b`, and master's `bf97700` (#522) landed after. Each branch has one commit the other lacks.
- ****There is no PR #525 / no NVENC / no `\_select\_encoder` anywhere in git history**** except that one unmerged branch. The `encoder` knob, `\_select\_encoder`, `\_probe\_encoder\_available`, and `docs/.../PHASE\_PLACEMENT.md` (whose §3 I'm told to update, and which already declares this exact "pre-rendered / overlay-filter watermark" follow-up) ****exist only on `feat/521-compile-on-l4`**** ? none are on master.
- So the task's literal instruction ("branch off `origin/master`") is incoherent as stated: master has neither the code this builds on nor the doc §3 to update. The 2026-07-08 CUJ §9 the task cites *is* on master (added by #522), and it confirms the target: stitch ran ****~12m55s**** (11 clips × ~67s/clip, watermark forcing a full 4K re-encode; concat only ~2s).

The work only makes sense on top of #521. That means stacking on an unmerged branch ? which your CLAUDE.md explicitly discourages ("NO PR stacking") ? so I want your call rather than picking it myself.

24. user (2026-07-08T06:00:13.933Z)

Your questions have been answered: "This task ("Part 3 of #521") builds on #521's encoder knob and updates PHASE_PLACEMENT.md §3 ? both of which exist **ONLY** on the unmerged, diverged branch origin/feat/521-compile-on-l4, not on origin/master. How should I base the work?"="The PR is merged so sync and build". You can now continue with these answers in mind.

25. assistant (2026-07-08T06:00:23.254Z)

Understood ? #521 merged on the remote after my initial fetch (local was lagging, exactly the parallel-slate scenario your CLAUDE.md warns about). Let me re-sync and confirm master now carries #521.

26. user (2026-07-08T06:00:31.807Z)

From <https://github.com/Khelsutra/badminton-highlight-indexer>

```
bf97700..dac6736 master -> origin/master
```

```
===origin/master now===
```

```
dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
```

```
placement (#521) (#525)
bf97700 docs(runlogs): record live SPA CUJ closing the 2 owner-gated validation-latency cells
($9) (#522)
699d85b feat(validation): cache validation verdict by video_id to skip re-validation on reupload
(#518)
```

```
===does origin/master contain #521 (8866ffe)?===
NO - #521 still not in master
```

```
===_select_encoder + PHASE_PLACEMENT on origin/master?===
PHASE_PLACEMENT.md EXISTS on master
_select_encoder count on master: 3
```

27. assistant (2026-07-08T06:04:14.362Z)

The user was right ? origin/master now advanced `bf97700 ? dac6736` = "perf(serving): run compile/stitch on the L4 worker ? (#521) (#525)". Master now has `PHASE\_PLACEMENT.md` and `\_select\_encoder` (count 3). My earlier fetch was simply stale. The dilemma is resolved: I branch off `origin/master`, which now literally contains #521.

Let me mark this chapter, set up tracking, and create the branch safely (I have unrelated dirty files ? the #523 handoff WIP and untracked `score\_\*.py` ? that I must not clobber or stage).

28. user (2026-07-08T06:04:15.808Z)

Chapter marked: "Ramp-up & base reconciliation". Continue your current work.

29. user (2026-07-08T06:04:20.135Z)

```
===does SESSION_HANDOFF.md differ HEAD vs origin/master?===
SAME (no conflict on checkout)
```

```
===ffmpeg + encoders===
ffmpeg version 8.1.1-full_build-www.gyan.dev Copyright (c) 2000-2026 the FFmpeg developers
V....D av1_nvenc          NVIDIA NVENC av1 encoder (codec av1)
V....D libx264            libx264 H.264 / AVC / MPEG-4 AVC / MPEG-4 part 10 (codec h264)
V....D libx264rgb         libx264 H.264 / AVC / MPEG-4 AVC / MPEG-4 part 10 RGB (codec h264)
V....D h264_nvenc         NVIDIA NVENC H.264 encoder (codec h264)
V....D libx265            libx265 H.265 / HEVC (codec hevc)
V....D hevc_nvenc         NVIDIA NVENC hevc encoder (codec hevc)
```

30. assistant (2026-07-08T06:04:29.965Z)

This box has `h264\_nvenc`/`hevc\_nvenc` *compiled in* (I'll probe runnability during benchmarking). Handoff won't conflict on checkout. Now I'll create the branch off `origin/master` per the branching-safety protocol, then verify HEAD is clean.

31. user (2026-07-08T06:04:33.738Z)

```
Exit code 1
error: Your local changes to the following files would be overwritten by checkout:
    backend/config/models.py
    backend/orchestration.py
    backend/pipeline/stitcher/compiler.py
    tests/test_compiler.py
Please commit your changes or stash them before you switch branches.
Aborting
```

32. assistant (2026-07-08T06:04:52.140Z)

Unexpected ? checkout is blocked by uncommitted local changes to `backend/config/models.py`, `orchestration.py`, `compiler.py`, `tests/test\_compiler.py`. These weren't in my session-start `git status` (which showed only `SESSION\_HANDOFF.md` + untracked `score\_\*.py`). This is the concurrent-shared-checkout hazard my CLAUDE.md warns about. Let me inspect carefully before

touching anything ? I must not clobber a sibling session's WIP.

33. user (2026-07-08T06:04:54.893Z)

===FULL git status===

On branch codex/issue-523-sd-reel-plan

Your branch is behind 'origin/master' by 1 commit, and can be fast-forwarded.

(use "git pull" to update your local branch)

Changes not staged for commit:

(use "git add <file>..." to update what will be committed)

(use "git restore <file>..." to discard changes in working directory)

modified: README.md
modified: SESSION_HANDOFF.md
modified: backend/config/models.py
modified: backend/infrastructure/database_jobs.py
modified: backend/main.py
modified: backend/orchestration.py
modified: backend/pipeline/stitcher/compiler.py
modified: config.json
modified: tests/test_api_endpoints.py
modified: tests/test_async_jobs.py
modified: tests/test_compiler.py
modified: tests/test_watermark.py

Untracked files:

(use "git add <file>..." to include in what will be committed)

score_golden_gate.py
score_local.py

no changes added to commit (use "git add" and/or "git commit -a")

===stat of the 4 flagged files vs HEAD (bf97700)===

warning: in the working copy of 'backend/config/models.py', LF will be replaced by CRLF the next time Git touches it

warning: in the working copy of 'backend/orchestration.py', LF will be replaced by CRLF the next time Git touches it

warning: in the working copy of 'backend/pipeline/stitcher/compiler.py', LF will be replaced by CRLF the next time Git touches it

warning: in the working copy of 'tests/test_compiler.py', LF will be replaced by CRLF the next time Git touches it

backend/config/models.py	24 ++++++++
backend/orchestration.py	30 ++++++++
backend/pipeline/stitcher/compiler.py	86 ++++++++-----
tests/test_compiler.py	64 ++++++++

4 files changed, 193 insertions(+), 11 deletions(-)

34. assistant (2026-07-08T06:05:52.584Z)

This is the concurrent-shared-checkout hazard my CLAUDE.md explicitly warns about. A **sibling session is actively implementing #523 (the SD-reel work)** in the shared tree right now ? the modified set (``main.py``, ``orchestration.py``, ``database_jobs.py``, ``compiler.py``, ``tests/test_compiler.py``, ``test_watermark.py``, ``config.json``, ``README.md``) matches the #523 ramp-up kit from the handoff exactly. That WIP includes ``compiler.py`` ? my target file. I must **not** clobber it, so I won't force the checkout.

The clean, non-destructive way to work concurrently is a **separate git worktree** off ``origin/master``. Let me verify the branch doesn't already exist and set one up in isolation.

35. user (2026-07-08T06:05:56.882Z)

===does branch exist?===

does not exist (good)

===existing worktrees===

```
C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer
bf97700 [codex/issue-523-sd-reel-plan]
C:/kg-wt521
8866ffe [feat/521-compile-on-l4]
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/cfg-wt      20b2d5a [claude/wasb-
config-validation]
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/compile-wt    b09a1d4 [claude/compile-
reliability]
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/pr458-pre-wt  355dad6 (detached HEAD)
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/pr458-wt      5de1ef2 (detached HEAD)
C:/Users/avidu/AppData/Local/Temp/claude/wt524-969
bf97700 [feat/524-l4-primary]
```

===create isolated worktree off origin/master===

```
Preparing worktree (new branch 'perf/521-cheap-watermark')
branch 'perf/521-cheap-watermark' set up to track 'origin/master'.
HEAD is now at dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven
phase?machine placement (#521) (#525)
```

36. assistant (2026-07-08T06:06:18.183Z)

Isolated worktree created at `C:/Users/avidu/Projects/khelsutra-guru/bhi-wt-cheap-watermark` on `perf/521-cheap-watermark` at `dac6736` (origin/master, includes #521). The worktree list confirms this is a heavily parallel environment (kg-wt521, wt524, compile-reliability, etc.) ? isolation was the right call. All my edits will happen in this worktree, leaving the sibling's #523 WIP untouched.

Let me confirm the worktree is clean and read the config surface + test patterns I'll build on.

37. user (2026-07-08T06:06:21.786Z)

===worktree status (should be clean)===
(clean if empty above)

===confirm #521 content present===

compiler _select_encoder refs: 5

===StitchingConfig in models.py===

```
37:         "recordings, and other encoders may use different GOP sizes ? tune "
332:class WatermarkConfig(BaseModel):
354:class StitchingConfig(BaseModel):
367:    # Video encoder for the per-clip trim re-encode (the branding watermark forces a re-
encode; the
369:    # ``hevc_nvenc`` are the L4 HARDWARE encoders ? dramatically faster on the 4K re-encode
that
370:    # dominated the 2026-07-08 CUJ stitch (#521). The compiler probes the requested encoder
ONCE and
372:    # cloud compile-dispatch forwards ``deployment.cloud_stitch_encoder`` onto this for the
L4 run;
374:    encoder: Literal["libx264", "h264_nvenc", "hevc_nvenc"] = "libx264"
477:class DeploymentConfig(BaseModel):
512:    # Video encoder the cloud compile-dispatch forwards onto ``stitching.encoder`` for the
L4 stitch
513:    # (#521). ``h264_nvenc`` uses the L4's hardware encoder for the 4K re-encode the
watermark forces;
515:    # the control-plane's own local stitch always uses ``stitching.encoder`` (libx264 by
default).
516:    cloud_stitch_encoder: Literal["libx264", "h264_nvenc", "hevc_nvenc"] = "h264_nvenc"
```

38. assistant (2026-07-08T06:06:25.935Z)

Worktree is clean with #521 content. Let me read the config classes and the test file to understand the surface I'll extend.

39. user (2026-07-08T06:06:27.042Z)

```
332     class WatermarkConfig(BaseModel):
333         """Branding overlay burned into each highlight clip during the per-clip re-encode
334         (backend/pipeline/stitcher/compiler.py). **On by default** ? set `enabled: false`
335         (under `stitching.watermark` in config.json) to turn it off. The text is fully
336         configurable. The concat stage stays stream-copy (-c copy)."""
337
338         enabled: bool = True
339         text: str = "Generated by Khelsutra.guru"
340         logo_path: str = "" # optional transparent PNG overlay; "" = no logo
341         font_path: str = "" # optional .ttf; "" = use the fontconfig `font` family below
342         font: str = "Sans" # fontconfig family used when font_path is empty
343         font_size_ratio: float = Field(
344             default=0.035, gt=0.0, le=0.5
345         ) # text height as a fraction of frame height
346         opacity: float = Field(default=0.85, ge=0.0, le=1.0)
347         box: bool = True # faint backing box behind the text for legibility
348         margin: int = Field(default=24, ge=0) # px inset from the chosen corner
349         position: Literal["bottom-right", "bottom-left", "top-right", "top-left"] = (
350             "bottom-right"
351         )
352
353
354     class StitchingConfig(BaseModel):
355         begin_padding: float = 0.5
356         end_padding: float = 0.5
357         use_cache: bool = False
358         min_shot_count: int = 1
359         # Long-rally reel filter (seconds; 0 = OFF). Drop detected rally pairs whose
duration
360         # (end ? start) is below this from the highlight reel, BEFORE compiling. A cleaner
361         # reel-selection knob than `min_shot_count`: shot_count is unlabeled / "Unknown" on
the
362         # fully-local no-AI path, whereas duration is derived from the rally's own
start/end. The
363         # local detector over-fires and its false positives are mostly SHORT (motion blips,
364         # background-court fragments), so this keeps the reel to the eye-catching long
rallies.
365         # OFF by default ? the detector output is unchanged, only which rallies reach the
reel.
366         min_rally_duration: float = Field(default=0.0, ge=0.0)
367         # Video encoder for the per-clip trim re-encode (the branding watermark forces a re-
encode; the
368         # concat stays -c copy). Default ``libx264`` (CPU, works everywhere incl. CI).
``h264_nvenc`` /
369         # ``hevc_nvenc`` are the L4 HARDWARE encoders ? dramatically faster on the 4K re-
encode that
370         # dominated the 2026-07-08 CUJ stitch (#521). The compiler probes the requested
encoder ONCE and
371         # falls back to libx264 if this ffmpeg build can't run it, so a non-GPU box never
breaks. The
372         # cloud compile-dispatch forwards ``deployment.cloud_stitch_encoder`` onto this for
the L4 run;
373         # the control-plane's own local stitch keeps libx264.
374         encoder: Literal["libx264", "h264_nvenc", "hevc_nvenc"] = "libx264"
375         watermark: WatermarkConfig = Field(default_factory=WatermarkConfig)
376
377
378     class LoggingConfig(BaseModel):
379         """Logging knobs for the run entrypoints (see backend/utils/logging_setup.py)."""
380
```

```

381     # Third-party HTTP/RPC client loggers (httpx, httpcore, urllib3, google-genai,
382     # google.auth, grpc) emit one INFO line per request ? dozens per Gemini run, which
383     # buries our own [rally]/[compile] lines in run.log during manual QA. Default raises
384     # them to WARNING; set true to restore full chatter for request-flow debugging.
385     verbose_http: bool = False
386
387
388     class SecurityConfig(BaseModel):
389         # When set, /api/process video_path must resolve under this directory (hosted/tenant
390         # mode). Default None preserves the single-user local 'any local file' workflow.
391         allowed_video_root: Optional[str] = None
392
393         # --- Chunked upload (B2, PR-2) ---
394         # Landing zone for browser uploads (assembled from parts). Per-user subdirectories
395         # arrive with PR-3 identity scoping. ? In hosted mode, set allowed_video_root to
396         # contain upload_dir or /api/process will reject the uploaded paths.
397         upload_dir: str = "output/uploads"
398         # Whole-file cap (default 4 GB) and per-part cap. Parts stay under ~100 MB because
399         # free-tier edge proxies cap request bodies there (HOSTING_PLAN $2).
400         max_upload_mb: int = 4096
401         max_part_mb: int = 95

```

40. user (2026-07-08T06:06:27.564Z)

```

1     """VideoCompiler coverage (previously zero ? the user-facing highlight-reel output path,
2     where the f-string SyntaxError lived undetected). ffmpeg/ffprobe are mocked."""
3
4     import os
5     import subprocess
6     from unittest.mock import MagicMock, patch
7
8     import pytest
9
10    from backend.pipeline.stitcher.compiler import VideoCompiler
11
12
13    def test_parse_time_formats():
14        f = VideoCompiler._parse_time
15        assert f(12.5) == 12.5
16        assert f("12.5") == 12.5
17        assert f("1:30") == 90.0
18        assert f("1:00:00") == 3600.0
19        assert f("garbage") == 0.0
20
21
22    def test_merge_overlapping_collapses_duplicates_and_overlaps():
23        m = VideoCompiler._merge_overlapping
24        # Exact duplicates + near-duplicates (the "every rally twice" re-run bug) collapse
25        # to one each.
26        assert m([(28.2, 34.5), (28.5, 34.5), (40.5, 45.5), (40.5, 45.5)]) == [
27            (28.2, 34.5),
28            (40.5, 45.5),
29        ]
30        # Unsorted input is sorted; an overlapping range extends the previous.
31        assert m([(100.0, 110.0), (0.0, 5.0), (104.0, 120.0)]) == [
32            (0.0, 5.0),
33            (100.0, 120.0),
34        ]
35        # Genuinely distinct (gapped) rallies are preserved, not merged.
36        assert m([(0.0, 5.0), (6.0, 9.0)]) == [(0.0, 5.0), (6.0, 9.0)]
37        # Degenerate rows (end <= start) are dropped; string timestamps are accepted.
38        assert m([("0:10", "0:20"), (30.0, 30.0), (31.0, 29.0)]) == [(10.0, 20.0)]
39
40    def test_compile_dedups_before_stitching(tmp_path):

```



```

41     """A segment list with every rally twice must yield ONE clip per unique rally."""
42     comp = VideoCompiler(str(tmp_path / "out"))
43     raw = tmp_path / "raw.mp4"
44     raw.write_bytes(b"x")
45     trims = []
46
47     def fake_run(cmd, *a, **k):
48         if cmd[0] == "ffprobe":
49             return MagicMock(stdout="100.0\n", returncode=0)
50         out = cmd[-1]
51         if isinstance(out, str) and out.endswith(".mp4"):
52             open(out, "wb").write(b"clip")
53             # count per-clip trims (the concat list file input is not an .mp4 output)
54             if "-i" in cmd and cmd.index("-i") + 1 == str(raw):
55                 trims.append(cmd)
56             return MagicMock(returncode=0, stdout=b"", stderr=b"")
57
58     dup_segments = [(0.0, 5.0), (0.2, 5.0), (10.0, 15.0), (10.0, 15.0), (20.0, 25.0)]
59     with patch.object(subprocess, "run", side_effect=fake_run):
60         comp.compile_highlights(str(raw), dup_segments, "reel.mp4")
61     # 5 input segments collapse to 3 unique rallies -> 3 per-clip trims, not 5.
62     assert len(trims) == 3, f"expected 3 deduped trims, got {len(trims)}"
63
64
65     def test_compile_basenames_output_and_stitches(tmp_path):
66         comp = VideoCompiler(str(tmp_path / "out"))
67         raw = tmp_path / "raw.mp4"
68         raw.write_bytes(b"x")
69
70         calls = []
71
72         def fake_run(cmd, *a, **k):
73             calls.append(cmd)
74             # Create whatever output file ffmpeg/ffprobe was told to write so existence
75             if cmd[0] == "ffprobe":
76                 return MagicMock(stdout="100.0\n", returncode=0)
77             out = cmd[-1]
78             if isinstance(out, str) and out.endswith(("mp4",)):
79                 open(out, "wb").write(b"clip")
80             return MagicMock(returncode=0, stdout=b"", stderr=b"")
81
82         with patch.object(subprocess, "run", side_effect=fake_run):
83             # Traversal in the name must be stripped (defense-in-depth basename).
84             out_path = comp.compile_highlights(
85                 str(raw), [(0.0, 5.0), (10.0, 15.0)], "../reel.mp4"
86             )
87
88             assert os.path.basename(out_path) == "reel.mp4"
89             assert os.path.dirname(out_path).endswith("out")
90             # Per-clip trims + one concat = at least 3 ffmpeg invocations (plus an ffprobe).
91             assert sum(1 for c in calls if c[0] == "ffmpeg") >= 3
92
93
94     def test_compile_concat_timeout_raises_clean_error(tmp_path):
95         comp = VideoCompiler(str(tmp_path / "out"))
96         raw = tmp_path / "raw.mp4"
97         raw.write_bytes(b"x")
98
99         def fake_run(cmd, *a, **k):
100             if cmd[0] == "ffprobe":
101                 return MagicMock(stdout="100.0\n", returncode=0)
102             if "concat" in cmd:
103                 raise subprocess.TimeoutExpired(cmd, k.get("timeout"))
104             out = cmd[-1]

```

```

105         if isinstance(out, str) and out.endswith(".mp4"):
106             open(out, "wb").write(b"clip")
107         return MagicMock(returncode=0, stdout=b"", stderr=b"")
108
109     with patch.object(subprocess, "run", side_effect=fake_run):
110         with pytest.raises(RuntimeError, match="timed out after 1800s"):
111             comp.compile_highlights(str(raw), [(0.0, 5.0)], "reel.mp4")
112
113
114 def test_concat_list_uses_absolute_clip_paths(tmp_path):
115     """Regression for Windows concat failure:
116
117     ffmpeg resolves relative paths inside clips.txt relative to clips.txt itself. If
118     clips.txt lives in output/tmp123 and contains output/tmp123/clip_0.mp4, ffmpeg
119     tries output/tmp123/output/tmp123/clip_0.mp4. The list must contain absolute
120     clip paths.
121     """
122     comp = VideoCompiler(str(tmp_path / "out"))
123     raw = tmp_path / "raw.mp4"
124     raw.write_bytes(b"x")
125     concat_list = ""
126
127     def fake_run(cmd, *a, **k):
128         nonlocal concat_list
129         if cmd[0] == "ffprobe":
130             return MagicMock(stdout="100.0\n", returncode=0)
131         if "-f" in cmd and "concat" in cmd:
132             list_path = cmd[cmd.index("-i") + 1]
133             concat_list = open(list_path, encoding="utf-8").read()
134         out = cmd[-1]
135         if isinstance(out, str) and out.endswith(".mp4"):
136             open(out, "wb").write(b"clip")
137         return MagicMock(returncode=0, stdout=b"", stderr=b"")
138
139     with patch.object(subprocess, "run", side_effect=fake_run):
140         comp.compile_highlights(str(raw), [(0.0, 5.0)], "reel.mp4")
141
142     line = concat_list.strip()
143     assert line.startswith("file ")
144     clip_path = line.removeprefix("file ").removesuffix("")
145     assert os.path.isabs(clip_path)
146     assert "clip_0.mp4" in clip_path
147
148
149 def test_compile_raises_on_no_segments(tmp_path):
150     comp = VideoCompiler(str(tmp_path / "out"))
151     raw = tmp_path / "raw.mp4"
152     raw.write_bytes(b"x")
153     with pytest.raises(ValueError):
154         comp.compile_highlights(str(raw), [], "reel.mp4")
155
156
157 def test_compile_raises_when_source_missing(tmp_path):
158     comp = VideoCompiler(str(tmp_path / "out"))
159     with pytest.raises(FileNotFoundError):
160         comp.compile_highlights(str(tmp_path / "nope.mp4"), [(0.0, 5.0)], "reel.mp4")
161
162
163 # ----- encoder selection (#521
? NVENC)
164
165
166 def test_select_encoder_default_is_libx264(tmp_path):
167     """Default encoder = libx264 (CPU) with the historical settings ? no probe, works
everywhere."""

```

```

168     comp = VideoCompiler(str(tmp_path / "out"))
169     args, chosen = comp._select_encoder()
170     assert chosen == "libx264"
171     assert args == ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"]
172
173
174     def test_select_encoder_nvenc_when_available(tmp_path, monkeypatch):
175         """h264_nvenc resolves to the hardware encoder args when the probe reports it
runnable (the L4)."""
176         from backend.pipeline.stitcher import compiler as comp_mod
177
178         monkeypatch.setattr(comp_mod, "_probe_encoder_available", lambda enc: True)
179         comp = VideoCompiler(str(tmp_path / "out"), encoder="h264_nvenc")
180         args, chosen = comp._select_encoder()
181         assert chosen == "h264_nvenc"
182         assert args[:2] == ["-c:v", "h264_nvenc"]
183         assert "-cq" in args # NVENC constant-quality knob (not -crf)
184
185
186     def test_select_encoder_falls_back_to_libx264_when_nvenc_unavailable(tmp_path,
monkeypatch):
187         """On a box where NVENC can't actually run (CI / no GPU), the requested encoder
falls back to
188         libx264 so a misconfig never nukes the reel."""
189         from backend.pipeline.stitcher import compiler as comp_mod
190
191         monkeypatch.setattr(comp_mod, "_probe_encoder_available", lambda enc: False)
192         comp = VideoCompiler(str(tmp_path / "out"), encoder="h264_nvenc")
193         args, chosen = comp._select_encoder()
194         assert chosen == "libx264"
195         assert args == ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"]
196
197
198     def test_compile_uses_selected_encoder_in_trim(tmp_path, monkeypatch):
199         """End-to-end (ffmpeg mocked): with an available NVENC encoder, every per-clip trim
command
200         carries -c:v h264_nvenc, and the concat stays -c copy."""
201         from backend.pipeline.stitcher import compiler as comp_mod
202
203         monkeypatch.setattr(comp_mod, "_probe_encoder_available", lambda enc: True)
204         comp = VideoCompiler(str(tmp_path / "out"), encoder="h264_nvenc")
205         raw = tmp_path / "raw.mp4"
206         raw.write_bytes(b"x")
207         trims = []
208
209         def fake_run(cmd, *a, **k):
210             if cmd[0] == "ffprobe":
211                 return MagicMock(stdout="100.0\n", returncode=0)
212             out = cmd[-1]
213             if isinstance(out, str) and out.endswith(".mp4"):
214                 open(out, "wb").write(b"clip")
215                 if "-i" in cmd and cmd.index("-i") + 1 == str(raw):
216                     trims.append(cmd)
217             return MagicMock(returncode=0, stdout=b"", stderr=b"")
218
219         with patch.object(subprocess, "run", side_effect=fake_run):
220             comp.compile_highlights(str(raw), [(0.0, 5.0), (10.0, 15.0)], "reel.mp4")
221         assert trims, "no trim commands captured"
222         for cmd in trims:
223             assert "h264_nvenc" in cmd # hardware encode on every clip
224             assert "libx264" not in cmd
225

```

I now understand the config surface (`StitchingConfig.encoder`, `WatermarkConfig`, `DeploymentConfig.cloud\_stitch\_encoder`) and the test patterns. Before implementing, the task says **measure on a real 4K clip** ? so I'll build an empirical benchmark to settle Option A vs B rather than theorize. First, probe whether NVENC actually **runs** on this box (it's compiled in), since the L4's default cloud encoder is `h264\_nvenc` ? that makes the benchmark far more representative.

42. user (2026-07-08T06:07:43.286Z)

```
===probe h264_nvenc runnable?===
h264_nvenc: NOT runnable (no GPU here)
hevc_nvenc: NOT runnable

===cpu count===
cpus: 16
===nvidia-smi?===
NVIDIA GeForce RTX 2070 Super with Max-Q Design
```

43. assistant (2026-07-08T06:07:56.985Z)

A GPU exists (RTX 2070 Super Max-Q) but NVENC failed to initialize. Let me see the actual error ? if it's a trivial probe issue I can use NVENC for a more L4-representative benchmark; if it's environmental, I'll benchmark on libx264 (the default anyway) and flag NVENC/L4 as owner-gated.

44. user (2026-07-08T06:08:00.167Z)

```
===h264_nvenc error (realistic probe)===
encoder      : Lavf62.12.101
Stream #0:0: Video: h264 (Main), yuv420p(tv, progressive), 1280x720 [SAR 1:1 DAR 16:9],
q=2-31, 2000 kb/s, 25 fps, 25 tbn
Metadata:
encoder      : Lavc62.28.101 h264_nvenc
Side data:
CPB properties: bitrate max/min/avg: 0/0/2000000 buffer size: 4000000 vbv_delay: N/A
[out#0/null @ 000001c452747b40] video:87KiB audio:0KiB subtitle:0KiB other streams:0KiB global
headers:0KiB muxing overhead: unknown
frame=      3 fps=0.0 q=22.0 Lsize=N/A time=00:00:00.00 bitrate=N/A speed=    0x elapsed=0:00:00.18
```

45. assistant (2026-07-08T06:08:21.350Z)

Interesting and important: `h264\_nvenc` **works** at 1280x720 ? it only failed at 64x64. That means NVENC **is** usable on this box (RTX 2070 Super), and I can make the benchmark representative of the L4's actual `h264\_nvenc` path. Let me confirm the exact #521 probe command's behavior vs a 4K encode.

46. user (2026-07-08T06:08:26.517Z)

```
===EXACT #521 probe (64x64)===
[out#0/null @ 0000020b1c847700] Nothing was written into output file, because at least one of
its streams received no packets.
frame=      0 fps=0.0 q=0.0 Lsize=      0KiB time=N/A bitrate=N/A speed=N/A elapsed=0:00:00.26
Conversion failed!
rc=0

===h264_nvenc at 4K (real use case)===
h264_nvenc 4K: OK
```

47. assistant (2026-07-08T06:09:29.465Z)

Confirmed two things:

- The exact #521 probe** (`color=...:s=64x64`) **fails** ? 64x64 is below NVENC's minimum encode size, so `\_probe\_encoder\_available` **under-reports** NVENC on this box (and potentially others). That's a latent bug in #521's probe, separate from my task ? I'll flag it, not fix it here.
- h264_nvenc works at 4K** ? so I can benchmark the L4's actual production encoder, not just

libx264.

Let me kick off generation of a realistic 4K 10-bit HEVC source (mimicking the corpus decode cost) in the background, and verify the bundled font while it runs.

48. user (2026-07-08T06:09:32.119Z)

```
total 1536
drwxr-xr-x 1 avidu 197609      0 Jul  8 11:35 .
drwxr-xr-x 1 avidu 197609      0 Jul  8 11:35 ..
-rw-r--r-- 2 avidu 197609 11560 Jul  8 11:35 LICENSE-Apache-2.0.txt
-rw-r--r-- 2 avidu 197609  4476 Jul  8 11:35 LICENSE-OFL-1.1.txt
-rw-r--r-- 2 avidu 197609 647144 Jul  8 11:35 NotoSansDevanagari-Variable.ttf
-rw-r--r-- 2 avidu 197609 641372 Jul  8 11:35 NotoSansKannada-Variable.ttf
-rw-r--r-- 2 avidu 197609  2100 Jul  8 11:35 README.md
-rw-r--r-- 2 avidu 197609 251880 Jul  8 11:35 RobotoSlab-Regular.ttf
---font exists check---
FONT OK
```

49. user (2026-07-08T06:09:37.353Z)

Command running in background with ID: bxcfpizw4. Output is being written to:
C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\tasks\bxcfpizw4.output. You will be notified when it completes. To check interim output, use Read on that file path.

50. assistant (2026-07-08T06:10:38.921Z)

Font present; source generating in the background. Let me write the benchmark harness while it generates. It mirrors the compiler's exact ffmpeg args and measures all four strategies (Baseline / Option A / Option B / no-watermark floor) across both encoders (libx264 default + h264_nvenc L4), with a per-stage (trim vs concat) breakdown ? the crux of whether cost actually *moves* under Option B.

51. assistant (2026-07-08T06:11:21.422Z)

"""Watermark-stitch cost benchmark for #521 Part-3 (cheap watermark).

Compares four stitch strategies on a real 4K clip, across encoders:

```
BASELINE : per-clip trim RE-ENCODE + drawtext watermark -> concat -c copy (current #521)
OPTION_A : pre-render watermark PNG once; per-clip trim RE-ENCODE + overlay -> concat -c
copy
OPTION_B : per-clip trim -c copy (stream-copy) -> concat RE-ENCODE + drawtext
watermark
FLOOR : per-clip trim -c copy -> concat -c copy (no watermark =
theoretical min)
```

Reports total wall + trim-stage vs concat-stage split, plus output validity (duration).

ffmpeg-only; mirrors backend/pipeline/stitcher/compiler.py arg construction.

"""

```
import os
import subprocess
import sys
import time
```

```
SP = os.path.dirname(os.path.abspath(__file__))
SRC = sys.argv[1] if len(sys.argv) > 1 else os.path.join(SP, "src_4k_hevc10.mp4")
FONT = r"C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-
watermark/backend/assets/fonts/RobotoSlab-Regular.ttf"
```

```
BEGIN_PAD, END_PAD = 0.5, 0.5
```

```
RAW_SEGMENTS = [
    (2.0, 6.5), (9.0, 13.5), (15.0, 19.0), (21.5, 25.0),
    (27.0, 31.5), (33.0, 37.0), (39.0, 43.0), (44.0, 47.0),
```

```

]

def ff_quote(path: str) -> str:
    return path.replace("\\", "/").replace("'", r"\'").replace(":", r"\:")

def probe_duration(path: str) -> float:
    try:
        out = subprocess.run(
            ["ffprobe", "-v", "error", "-show_entries", "format=duration",
             "-of", "default=noprint_wrappers=1:nokey=1", path],
            capture_output=True, text=True, timeout=60,
        )
        return float(out.stdout.strip())
    except Exception:
        return 0.0

def probe_res(path: str):
    out = subprocess.run(
        ["ffprobe", "-v", "error", "-select_streams", "v:0", "-show_entries",
         "stream=width,height", "-of", "csv=p=0:s=x", path],
        capture_output=True, text=True, timeout=60,
    )
    w, h = out.stdout.strip().split("x")
    return int(w), int(h)

def drawtext_filter(tmp: str) -> str:
    tf = os.path.join(tmp, "watermark.txt")
    with open(tf, "w", encoding="utf-8") as fh:
        fh.write("Generated by Khelsutra.guru")
    # divisor = round(1/0.035) = 29 ; matches compiler _watermark_filter
    return (
        f"drawtext=textfile='{ff_quote(tf)}':fontfile='{ff_quote(FONT)}'"
        f":fontcolor=white@0.85:fontsize=h/29:x=w-tw-24:y=h-th-24"
        f":box=1:boxcolor=black@0.50:boxborderw=10"
    )

def enc_args(encoder: str):
    if encoder == "libx264":
        return ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"]
    return ["-c:v", encoder, "-preset", "p4", "-cq", "23", "-pix_fmt", "yuv420p"]

def run(cmd) -> float:
    t = time.perf_counter()
    p = subprocess.run(cmd, capture_output=True)
    dt = time.perf_counter() - t
    if p.returncode != 0:
        raise RuntimeError("ffmpeg failed:\n" + " ".join(str(c) for c in cmd)
                           + "\n" + p.stderr.decode(errors="replace")[:-800:])
    return dt

def padded(seg):
    s, e = seg
    return max(0.0, s - BEGIN_PAD), e + END_PAD

def clips_txt(tmp, clips):
    lp = os.path.join(tmp, "list.txt")
    with open(lp, "w", encoding="utf-8") as f:

```

```

        for c in clips:
            f.write("file '%s'\n" % c.replace("\\", "/").replace("'", "'\\'"))
    return lp

def do_baseline(tmp, enc):
    va = enc_args(enc)
    dt = drawtext_filter(tmp)
    trim_t = 0.0
    clips = []
    for i, seg in enumerate(RAW_SEGMENTS):
        ps, pe = padded(seg)
        out = os.path.join(tmp, f"b_{i}.mp4")
        trim_t += run(["ffmpeg", "-y", "-ss", str(ps), "-to", str(pe), "-i", SRC,
                        "-vf", dt, *va, "-c:a", "aac", "-b:a", "128k", out])
        clips.append(out)
    lp = clips_txt(tmp, clips)
    outp = os.path.join(tmp, "reel_baseline.mp4")
    concat_t = run(["ffmpeg", "-y", "-f", "concat", "-safe", "0", "-i", lp, "-c", "copy", outp])
    return trim_t, concat_t, outp

def do_option_a(tmp, enc):
    va = enc_args(enc)
    dt = drawtext_filter(tmp)
    W, H = probe_res(SRC)
    wm = os.path.join(tmp, "wm.png")
    # Pre-render the SAME drawtext onto a transparent RGBA canvas ONCE.
    prerender_t = run(["ffmpeg", "-y", "-f", "lavfi", "-i",
                        f"color=c=black@0.0:s={W}x{H}:r=1:d=1,format=rgba",
                        "-vf", dt, "-frames:v", "1", wm])
    ov = "W-w-24:H-h-24"
    wmq = ff_quote(wm)
    trim_t = 0.0
    clips = []
    for i, seg in enumerate(RAW_SEGMENTS):
        ps, pe = padded(seg)
        out = os.path.join(tmp, f"a_{i}.mp4")
        trim_t += run(["ffmpeg", "-y", "-ss", str(ps), "-to", str(pe), "-i", SRC,
                        "-vf", f"movie='{wmq}'[wm];[in][wm]overlay={ov}",
                        *va, "-c:a", "aac", "-b:a", "128k", out])
        clips.append(out)
    lp = clips_txt(tmp, clips)
    outp = os.path.join(tmp, "reel_optA.mp4")
    concat_t = run(["ffmpeg", "-y", "-f", "concat", "-safe", "0", "-i", lp, "-c", "copy", outp])
    return trim_t, concat_t, outp, prerender_t

def do_option_b(tmp, enc):
    va = enc_args(enc)
    dt = drawtext_filter(tmp)
    trim_t = 0.0
    clips = []
    for i, seg in enumerate(RAW_SEGMENTS):
        ps, pe = padded(seg)
        dur = pe - ps
        out = os.path.join(tmp, f"c_{i}.mp4")
        # stream-copy trim: fast seek to keyframe <= ps, copy dur seconds (frame-INACCURATE at
start).
        trim_t += run(["ffmpeg", "-y", "-ss", str(ps), "-i", SRC, "-t", str(dur),
                        "-c", "copy", "-avoid_negative_ts", "make_zero", out])
        clips.append(out)
    lp = clips_txt(tmp, clips)
    outp = os.path.join(tmp, "reel_optB.mp4")
    # single re-encode pass over the concatenated reel, applying the watermark once.

```

[illegible]


```

print("\n\n==== SUMMARY (measured wall seconds) ====")
print(f"{'encoder':12s} {'strategy':10s} {'total':>8s} {'trim':>8s} {'concat':>8s}
{'reel_s':>7s} {'MB':>7s}")
for enc, name, total, trim_t, concat_t, dur, sz, err in rows:
    if total is None:
        print(f"{enc:12s} {name:10s} ERROR: {err}")
    else:
        print(f"{enc:12s} {name:10s} {total:8.2f} {trim_t:8.2f} {concat_t:8.2f} {dur:7.1f}
{sz:7.1f}")

if __name__ == "__main__":
    main()

```

52. user (2026-07-08T06:11:21.470Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\scratchpad\wm_bench.py (file state is current in your context ? no need to Read it back)

53. assistant (2026-07-08T06:11:26.400Z)

Let me check whether the source finished generating.

54. user (2026-07-08T06:11:30.766Z)

```

-rw-r--r-- 1 avidu 197609 65273900 Jul  8 11:40 C:/Users/avidu/AppData/Local/Temp/claude/C--
Users-avidu-Projects-khelsutra-guru-badminton-highlight-
indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/src_4k_hevc10.mp4
---tail of gen log---
Generating 4K 10-bit HEVC source (mandelbrot, detailed/realistic decode cost, GOP 60, ~48s)...

```

55. assistant (2026-07-08T06:13:03.862Z)

While the source generates, let me reconsider the frame-accuracy question, since it's decisive.

****Key insight:**** the watermark burns into *every* frame, so every frame must be re-encoded regardless. That reframes both options:

- ****Option A**** (pre-rendered overlay, per-clip): frame-accurate (re-encode allows precise ``-ss`` cuts), same encode volume as baseline, saves only per-frame ``drawtext`` rasterization.
- ****Option B**** (stream-copy trims + concat re-encode): the concat **still** re-encodes the whole reel for the watermark, so the encode floor doesn't shrink ? it just **moves**. And stream-copy trims start at keyframes, so each clip is keyframe-extended (encodes **more** footage) **and** loses frame-accuracy. The "smart-cut boundary GOP" hybrid the task mentions is moot here ? you're re-encoding every frame for the watermark anyway.

So theory predicts Option B is a wash-or-worse when the watermark is on, and the real per-clip cost is the irreducible encode. The benchmark's ``FLOOR`` (stream-copy only) vs ``BASELINE`` decomposition will confirm how much is overhead vs encode. Let me check the source.

56. user (2026-07-08T06:13:08.601Z)

```

-rw-r--r-- 1 avidu 197609 122945580 Jul  8 11:43 C:/Users/avidu/AppData/Local/Temp/claude/C--
Users-avidu-Projects-khelsutra-guru-badminton-highlight-
indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/src_4k_hevc10.mp4
---
[mov,mp4,m4a,3gp,3g2,mj2 @ 0000025490481700] moov atom not found
C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-
highlight-indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/src_4k_hevc10.mp4: Invalid
data found when processing input
still writing / not finalized

```

57. assistant (2026-07-08T06:13:36.218Z)

Source still finalizing (mandelbrot 4K is CPU-heavy; ~122MB of ~300MB). I'll wait for the completion notification. Meanwhile, let me study the approach-independent wiring ? how `VideoCompiler` is constructed from config in orchestration, so I know exactly where a new knob threads through.

58. user (2026-07-08T06:13:39.651Z)

```
===VideoCompiler construction sites===
backend/main.py:1381:         compiler = VideoCompiler(
backend/orchestration.py:1215:     compiler = _main.VideoCompiler(

===_stitch_reel / compile core in orchestration===
16: ``ProcessRequest``, ``compute_video_id``, ``emit_indexer_report``, ``VideoCompiler``.
1105:     and survive the ``stitching.min_shot_count`` floor (segments WITHOUT a known shot_count
are exempt
1128:         if sc is not None and int(sc) < stitching.min_shot_count:
1135:         "[compile] stitching.min_shot_count=%s: dropped %d/%d segments below the
floor.",
1136:         stitching.min_shot_count,
1144:         f"stitching.min_shot_count={stitching.min_shot_count}.".
1190: def _stitch_reel(
1202:     needs NO control-plane DB ? the worker reads them from the staged spec. Builds the
VideoCompiler
1203:     with ``stitching.encoder`` (NVENC on the L4, libx264 on the CP), fetches the source,
runs the
1207:     import backend.main as _main # call-time seam: VideoCompiler stays patchable on
backend.main
1214:     localized_watermark = localize_watermark(stitching.watermark, locale)
1215:     compiler = _main.VideoCompiler(
1217:         begin_padding=stitching.begin_padding,
1218:         end_padding=stitching.end_padding,
1219:         watermark=localized_watermark,
1220:         encoder=getattr(stitching, "encoder", "libx264"),
1326:     stitch_dump = stitching.model_dump()
```

59. assistant (2026-07-08T06:13:44.251Z)

Let me read both `VideoCompiler` construction sites to see the exact config threading.

60. user (2026-07-08T06:13:45.774Z)

```
1190     def _stitch_reel(
1191         cfg,
1192         stitching,
1193         video_id: str,
1194         source_ref: str,
1195         time_pairs: List[Tuple[float, float]],
1196         out_name: str,
1197         locale: Optional[str],
1198         notify,
1199     ) -> Dict[str, Any]:
1200         """Stitch + mirror-upload core, shared by the control-plane compile handler AND the
L4 worker
1201         (#521). ``time_pairs`` are already RESOLVED (the caller filtered via
``_resolve_compile``), so this
1202         needs NO control-plane DB ? the worker reads them from the staged spec. Builds the
VideoCompiler
1203         with ``stitching.encoder`` (NVENC on the L4, libx264 on the CP), fetches the source,
runs the
1204         trim/watermark/concat, mirrors the reel to
``<output_root>/highlights/<video_id>/<name>``, and
1205         returns the ``{status, filepath, download_url, video_id, reel_uri}`` the SPA polls
for. Returns a
1206         ``status='failed'`` result on a stitch error (never raises for an expected ffmpeg
```

```

failure)."""
1207     import backend.main as _main # call-time seam: VideoCompiler stays patchable on
backend.main
1208
1209     output_dir = (
1210         getattr(getattr(cfg, "output", None), "output_dir", "output") or "output"
1211     )
1212     # Localize the branding watermark to the UI locale the SPA rendered in (recorded on
the job);
1213     # unknown/absent locale keeps the configured default unchanged.
1214     localized_watermark = localize_watermark(stitching.watermark, locale)
1215     compiler = _main.VideoCompiler(
1216         output_dir,
1217         begin_padding=stitching.begin_padding,
1218         end_padding=stitching.end_padding,
1219         watermark=localized_watermark,
1220         encoder=getattr(stitching, "encoder", "libx264"),
1221     )
1222
1223     notify("stitching")
1224     local_src = ""
1225     try:
1226         # The stored `filepath` is the SOURCE ref (gs://, s3://, gdrive://, or a local
path). Resolve to
1227         # a LOCAL file the stitcher can read ? identity for a local path, fetch-to-
scratch for a URI (the
1228         # same seam the process path uses); without it, compiling a bucket-ingested
video can't open gs://.
1229         local_src = storage.acquire_local_input(
1230             source_ref, getattr(cfg, "storage", None)
1231         )
1232         # The ffmpeg stitch is CPU-heavy ? serialize it so a raised drain worker count
(for remote
1233         # dispatch) doesn't run multiple compiles at once on the box (see
_LOCAL_COMPUTE_LANE).
1234         with _LOCAL_COMPUTE_LANE:
1235             out_path = compiler.compile_highlights(local_src, time_pairs, out_name)
1236     except Exception as e: # noqa: BLE001 ? a stitch/ffmpeg failure is a clean job
FAILURE, not a crash
1237         logger.error(
1238             "[compile] stitching failed for video_id=%s: %s", video_id, e, exc_info=True
1239         )

```

61. user (2026-07-08T06:13:46.328Z)

```

1370         out_name = out_filename
1371
1372         os.makedirs(out_dir, exist_ok=True)
1373         print(
1374             f"[CLI] Extracting debug segment from {args.trim_start}s to {args.trim_end}s
from {args.video_path}..."
1375         )
1376
1377         # global_config isn't initialized this early; load the typed config so the trim
1378         # clip carries the same configured stitching paddings + watermark as a real
1379         # highlight (watermark default-on).
1380         stitching = load_app_config().stitching
1381         compiler = VideoCompiler(
1382             out_dir,
1383             begin_padding=stitching.begin_padding,
1384             end_padding=stitching.end_padding,
1385             watermark=stitching.watermark,
1386         )
1387         try:
1388             res_path = compiler.compile_highlights(

```

```

1389         args.video_path, [(args.trim_start, args.trim_end)], out_name
1390     )
1391     print(f"[SUCCESS] Trimmed segment saved to: {res_path}")
1392 except Exception as e:
1393     print(f"[ERROR] Failed to compile trimmed segment: {e}")
1394     return
1395
1396 if args.debug_clip is not None:
1397     if not args.video_path:
1398         print("[ERROR] Please provide --video-path to debug a specific clip.")
1399     return

```

62. assistant (2026-07-08T06:14:00.096Z)

Full config flow understood: a new knob on `StitchingConfig` threads through `\_stitch\_reel` ? `VideoCompiler` (and rides the cloud compile-spec automatically since the spec carries `stitching`). Same pattern as `encoder`. Let me check the source generation.

63. user (2026-07-08T06:14:05.060Z)

```

-rw-r--r-- 1 avidu 197609 162529324 Jul  8 11:44 C:/Users/avidu/AppData/Local/Temp/claude/C--
Users-avidu-Projects-khelsutra-guru-badminton-highlight-
indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/src_4k_hevc10.mp4
still generating

```

64. assistant (2026-07-08T06:14:25.086Z)

162MB and climbing. Rather than poll, let me do genuinely independent work: rigorously characterize the NVENC-probe bug I found (to make the flag actionable) by testing which resolutions the probe passes/fails ? this pins down the fix.

65. user (2026-07-08T06:14:33.424Z)

```

color 64x64: FAIL
color 128x128: FAIL
color 145x49: PASS
color 256x256: PASS
color 320x240: PASS
--- with a real (testsrc2) source, same tiny sizes ---
testsrc2 64x64: FAIL
testsrc2 128x128: FAIL
testsrc2 256x256: PASS

```

66. assistant (2026-07-08T06:14:59.413Z)

This precisely characterizes a **latent bug in #521**: `\_probe\_encoder\_available` uses a `64x64` probe, but NVENC's minimum encode size is higher ? `64x64`/`128x128` **fail** while `145x49`/`256x256` **pass**, even with a real source. So on any box with a working NVENC (including potentially the L4), the probe **under-reports** NVENC as unavailable and silently falls back to libx264 ? defeating #521's core optimization. It's out of scope for my watermark PR, but worth flagging with the concrete fix while it's fresh.

67. user (2026-07-08T06:15:10.730Z)

Noted (position 1, task_id: task_8a2b7d23). A chip is showing for the user ? they can start it in a fresh worktree with one click, or dismiss it. If this suggestion becomes stale or superseded, call `dismiss_task` with this task_id. Currently pending: task_8a2b7d23 "Fix NVENC probe: 64x64 under-reports availability". Continue your current work.

68. assistant (2026-07-08T06:15:15.848Z)

Flagged. Let me check the source status.

69. user (2026-07-08T06:15:20.754Z)

```
-rw-r--r-- 1 avidu 197609 232521772 Jul  8 11:45 C:/Users/avidu/AppData/Local/Temp/claude/C--
Users-avidu-Projects-khelsutra-guru-badminton-highlight-
indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/src_4k_hevc10.mp4
still generating
```

70. assistant (2026-07-08T06:15:33.434Z)

232MB ? nearly done. Rather than poll, let me use the wait productively: read the existing watermark tests (clean origin/master version in my worktree) to understand the test surface I'll extend regardless of which option wins.

71. user (2026-07-08T06:15:34.156Z)

```
1      """Tests for the configurable highlight watermark (stitching.watermark).
2
3      Covers the config defaults, the filtergraph builder (text + logo, positions, font,
4      the configurable string), and the end-to-end wiring through compile_highlights with
5      ffmpeg/ffprobe mocked ? including the graceful fallback when the watermark encode fails.
6      """
7
8      import os
9      import subprocess
10     from pathlib import Path
11     from unittest.mock import MagicMock, patch
12
13     import pytest
14     from pydantic import ValidationError
15
16     from backend.config import AppConfig
17     from backend.config.models import WatermarkConfig
18     from backend.pipeline.stitcher.compiler import VideoCompiler
19
20     # --- config -----
21
22
23     def test_watermark_config_defaults_on_with_brand_text():
24         wm = AppConfig().stitching.watermark
25         assert wm.enabled is True # default-on, with an off switch
26         assert wm.text == "Generated by Khelsutra.guru"
27         assert wm.position == "bottom-right"
28
29
30     def test_watermark_can_be_disabled_via_config():
31         cfg = AppConfig.model_validate({"stitching": {"watermark": {"enabled": False}}})
32         assert cfg.stitching.watermark.enabled is False
33
34
35     def test_font_size_ratio_and_opacity_bounds_enforced():
36         with pytest.raises(ValidationError):
37             WatermarkConfig(opacity=1.5)
38         with pytest.raises(ValidationError):
39             WatermarkConfig(font_size_ratio=0.0)
40
41
42     # --- filtergraph builder -----
43
44
45     def test_disabled_or_none_builds_no_filter(tmp_path):
46         assert (
47             VideoCompiler(str(tmp_path / "a"), watermark=None)._watermark_filter(
48                 str(tmp_path)
49             )
50             is None
```

```

51         )
52         off = VideoCompiler(str(tmp_path / "b"), watermark={"enabled": False})
53         assert off._watermark_filter(str(tmp_path)) is None
54
55
56     def test_text_filter_writes_configurable_string_and_builds_drawtext(tmp_path):
57         custom = "My Academy 2026 ? Khelsutra.guru"
58         comp = VideoCompiler(
59             str(tmp_path / "o"), watermark={"enabled": True, "text": custom}
60         )
61         graph = comp._watermark_filter(str(tmp_path))
62         assert graph is not None and graph.startswith("drawtext=")
63         # The arbitrary, configurable string is written verbatim to the textfile (no
64         # escaping pain).
65         assert (tmp_path / "watermark.txt").read_text(encoding="utf-8") == custom
66         assert "textfile=" in graph
67         assert "fontcolor=white@0.85" in graph
68         assert "fontsize=h/29" in graph # round(1/0.035) == 29
69         assert "x=w-tw-24:y=h-th-24" in graph # bottom-right default, margin 24
70         assert "box=1" in graph
71
72     def test_position_and_margin_map_to_corner(tmp_path):
73         comp = VideoCompiler(
74             str(tmp_path / "o"),
75             watermark={"enabled": True, "text": "X", "position": "top-left", "margin": 12},
76         )
77         graph = comp._watermark_filter(str(tmp_path))
78         assert "x=12:y=12" in graph
79
80
81     def test_explicit_font_path_uses_fontfile(tmp_path):
82         comp = VideoCompiler(
83             str(tmp_path / "o1"),
84             watermark={"enabled": True, "text": "X", "font_path": "/fonts/Brand.ttf"},
85         )
86         g = comp._watermark_filter(str(tmp_path))
87         assert "fontfile='/fonts/Brand.ttf'" in g
88
89
90     def test_default_uses_bundled_apache_font(tmp_path):
91         # No font_path ? the bundled Apache-2.0 Roboto Slab (renders without fontconfig).
92         comp = VideoCompiler(str(tmp_path / "o2"), watermark={"enabled": True, "text": "X"})
93         g = comp._watermark_filter(str(tmp_path))
94         assert "fontfile=" in g and "RobotoSlab-Regular.ttf" in g
95         assert "font=" not in g # not the fontconfig path
96
97
98     def test_bundled_font_actually_ships_and_is_apache():
99         import backend.pipeline.stitcher.compiler as cmod
100
101         assert cmod._BUNDLED_FONT.exists(), (
102             "bundled default watermark font must ship in the repo"
103         )
104         with open(cmod._BUNDLED_FONT, "rb") as fh:
105             assert fh.read(4) in (
106                 b"\x00\x01\x00\x00",
107                 b"OTTO",
108                 b"true",
109                 b"ttcf",
110             ) # valid sfnt
111         lic = cmod._BUNDLED_FONT.parent / "LICENSE-Apache-2.0.txt"
112         assert lic.exists() and "Apache License" in lic.read_text(
113             encoding="utf-8", errors="replace"
114         )

```

```

115
116
117 def test_fontconfig_fallback_when_bundled_missing(tmp_path, monkeypatch):
118     import backend.pipeline.stitcher.compiler as cmod
119
120     monkeypatch.setattr(cmod, "_BUNDLED_FONT", Path("/no/such/font.ttf"))
121     comp = VideoCompiler(
122         str(tmp_path / "o3"),
123         watermark={"enabled": True, "text": "X", "font": "DejaVu Sans"},
124     )
125     g = comp._watermark_filter(str(tmp_path))
126     assert "font='DejaVu Sans'" in g and "fontfile=" not in g
127
128
129 # --- Windows drive-letter colon escaping (regression) -----
130 # A Windows absolute path's drive-letter colon (C:/...) must be escaped for the
131 # ffmpeg filtergraph, or drawtext/movie parse ':' as an option separator and the
132 # encode fails with "Invalid argument" ? which silently disabled the on-by-default
133 # bundled-font watermark (and absolute logo paths) on Windows. Verified end-to-end
134 # against real ffmpeg before this fix landed.
135
136
137 def test_ff_quote_escapes_drive_letter_colon():
138     assert VideoCompiler._ff_quote(r"C:\fonts\Brand.ttf") == r"C\:/fonts/Brand.ttf"
139     # no colon -> unchanged (POSIX paths and the Linux bundled font are unaffected)
140     assert (
141         VideoCompiler._ff_quote("/usr/share/fonts/font.ttf")
142         == "/usr/share/fonts/font.ttf"
143     )
144     # colon + single-quote + backslash all handled
145     assert VideoCompiler._ff_quote(r"D:\a\b\f.ttf") == r"D\:/a\b/f.ttf"
146
147
148 def test_watermark_font_path_with_drive_colon_is_escaped(tmp_path):
149     # Explicit font_path (literal, existence not required) -> deterministic on any OS.
150     comp = VideoCompiler(
151         str(tmp_path / "o"),
152         watermark={"enabled": True, "text": "X", "font_path": "C:/fonts/Brand.ttf"},
153     )
154     g = comp._watermark_filter(str(tmp_path))
155     assert r"fontfile='C\:/fonts/Brand.ttf'" in g
156     assert "fontfile='C:/fonts'" not in g # the raw colon would break drawtext parsing
157
158
159 def test_logo_plus_text_overlays_then_draws(tmp_path):
160     logo = tmp_path / "logo.png"
161     logo.write_bytes(b"\x89PNG\r\n")
162     comp = VideoCompiler(
163         str(tmp_path / "o"),
164         watermark={"enabled": True, "text": "X", "logo_path": str(logo)},
165     )
166     graph = comp._watermark_filter(str(tmp_path))
167     assert graph.startswith("movie=")
168     assert "overlay=W-w-24:H-h-24" in graph
169     assert "drawtext=" in graph
170
171
172 def test_missing_logo_is_skipped_gracefully(tmp_path):
173     comp = VideoCompiler(
174         str(tmp_path / "o"),
175         watermark={
176             "enabled": True,
177             "text": "X",
178             "logo_path": str(tmp_path / "nope.png"),
179         },

```

```

180     )
181     graph = comp._watermark_filter(str(tmp_path))
182     assert "movie=" not in graph and graph.startswith("drawtext=") # text still drawn
183
184
185     # --- end-to-end wiring (ffmpeg/ffprobe mocked) -----
186
187
188     def _fake_run_factory(calls, fail_on_vf=False):
189         def fake_run(cmd, *a, **k):
190             calls.append(cmd)
191             if cmd[0] == "ffprobe":
192                 return MagicMock(stdout="100.0\n", returncode=0)
193             if fail_on_vf and "-vf" in cmd:
194                 raise subprocess.CalledProcessError(
195                     1, cmd, output=b"", stderr=b"font not found"
196                 )
197             out = cmd[-1]
198             if isinstance(out, str) and out.endswith(".mp4"):
199                 open(out, "wb").write(b"clip")
200             return MagicMock(returncode=0, stdout=b"", stderr=b"")
201
202         return fake_run
203
204
205     def test_compile_applies_watermark_vf_to_each_clip_not_concat(tmp_path):
206         raw = tmp_path / "raw.mp4"
207         raw.write_bytes(b"x")
208         comp = VideoCompiler(
209             str(tmp_path / "out"),
210             watermark={"enabled": True, "text": "Generated by Khelsutra.guru"},
211         )
212         calls = []
213         with patch.object(subprocess, "run", side_effect=_fake_run_factory(calls)):
214             comp.compile_highlights(str(raw), [(0.0, 5.0), (10.0, 15.0)], "reel.mp4")
215
216         trim_cmds = [
217             c for c in calls if c[0] == "ffmpeg" and "-i" in c and "concat" not in c
218         ]
219         assert trim_cmds, "expected per-clip trim invocations"
220         for c in trim_cmds:
221             assert "-vf" in c
222             assert "drawtext=" in c[c.index("-vf") + 1]
223         # The concat (-c copy) must NOT re-filter ? the mark is already baked into each
clip.
224         concat_cmds = [c for c in calls if "concat" in c]
225         assert concat_cmds and all("-vf" not in c for c in concat_cmds)
226
227
228     def test_compile_without_watermark_adds_no_vf(tmp_path):
229         raw = tmp_path / "raw.mp4"
230         raw.write_bytes(b"x")
231         comp = VideoCompiler(str(tmp_path / "out")) # no watermark passed
232         calls = []
233         with patch.object(subprocess, "run", side_effect=_fake_run_factory(calls)):
234             comp.compile_highlights(str(raw), [(0.0, 5.0)], "reel.mp4")
235         assert all("-vf" not in c for c in calls)
236
237
238     def test_watermark_failure_falls_back_to_unwatermarked_clip(tmp_path):
239         """A bad font/logo must not nuke the reel: the clip re-encodes without the
watermark."""
240         raw = tmp_path / "raw.mp4"
241         raw.write_bytes(b"x")
242         comp = VideoCompiler(

```



```

243         str(tmp_path / "out"),
244         watermark={"enabled": True, "text": "Generated by Khelsutra.guru"},
245     )
246     calls = []
247     with patch.object(
248         subprocess, "run", side_effect=_fake_run_factory(calls, fail_on_vf=True)
249     ):
250         out_path = comp.compile_highlights(str(raw), [(0.0, 5.0)], "reel.mp4")
251
252     # Reel still produced despite the watermark encode failing.
253     assert out_path.endswith("reel.mp4")
254     assert os.path.exists(out_path)
255     # For the clip, both a watermarked attempt and a plain fallback attempt were made.
256     clip_cmds = [
257         c for c in calls if c[0] == "ffmpeg" and "concat" not in c and "-i" in c
258     ]
259     assert any("-vf" in c for c in clip_cmds) # watermark attempt
260     assert any("-vf" not in c for c in clip_cmds) # fallback attempt
261
262
263 def test_watermark_failure_logs_the_ffmpeg_reason(tmp_path, caplog):
264     """The fallback warning must surface ffmpeg's stderr (the *why*), not go silent ?
265     e.g. the 'Fontconfig error' that tells the user to set font_path."""
266     import logging as _logging
267
268     raw = tmp_path / "raw.mp4"
269     raw.write_bytes(b"x")
270     comp = VideoCompiler(
271         str(tmp_path / "out"), watermark={"enabled": True, "text": "Y"}
272     )
273
274     def fake_run(cmd, *a, **k):
275         if cmd[0] == "ffprobe":
276             return MagicMock(stdout="100.0\n", returncode=0)
277         if "-vf" in cmd:
278             raise subprocess.CalledProcessError(
279                 1,
280                 cmd,
281                 output=b"",
282                 stderr=b"Fontconfig error: Cannot load default config file",
283             )
284         out = cmd[-1]
285         if isinstance(out, str) and out.endswith(".mp4"):
286             open(out, "wb").write(b"clip")
287         return MagicMock(returncode=0, stdout=b"", stderr=b"")
288
289     with caplog.at_level(_logging.WARNING):
290         with patch.object(subprocess, "run", side_effect=fake_run):
291             comp.compile_highlights(str(raw), [(0.0, 5.0)], "reel.mp4")
292     assert "Fontconfig error" in caplog.text
293

```

72. assistant (2026-07-08T06:15:43.367Z)

I now have the complete test surface. Critically, ``test_compile_applies_watermark_vf_to_each_clip_not_concat`` locks the **current** contract (watermark per-clip, concat ``-c copy``, no ``-vf`` on concat) ? so any Option-B behavior must be ***opt-in behind a knob*** with the default keeping that test green. Let me check the source.

73. user (2026-07-08T06:15:48.259Z)

```

-rw-r--r-- 1 avidu 197609 309665754 Jul  8 11:45 C:/Users/avidu/AppData/Local/Temp/claude/C--
Users-avidu-Projects-khelsutra-guru-badminton-highlight-
indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/src_4k_hevc10.mp4
codec_name=hevc

```

```
width=3840
height=2160
pix_fmt=yuv420p10le
duration=48.000000
FINALIZED
```

74. assistant (2026-07-08T06:16:06.534Z)

Source finalized: **3840×2160, yuv420p10le (10-bit HEVC), 48s, ~310MB (~52 Mbps)** ? a realistically decode-heavy 4K clip mirroring the corpus. Now let me run the benchmark (Baseline/A/B/Floor × libx264+NVENC, warm+measured) in the background and monitor.

75. user (2026-07-08T06:16:08.373Z)

Command running in background with ID: br1xehp9b. Output is being written to: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\tasks\br1xehp9b.output. You will be notified when it completes. To check interim output, use Read on that file path.

76. assistant (2026-07-08T06:16:19.674Z)

Benchmark running. Let me check interim output to confirm it started cleanly (encoders detected, no immediate errors).

77. user (2026-07-08T06:16:45.387Z)

(Bash completed with no output)

78. assistant (2026-07-08T06:17:30.883Z)

Python buffers stdout when redirected, so interim prints won't flush until exit ? I'll get the full table on completion. Let me verify the benchmark is actively progressing (not wedged) by checking for the temp clip files it creates, rather than waiting blind.

79. user (2026-07-08T06:17:45.356Z)

```
===python running?===
python.exe                11352 Console                2      14,200 K
python.exe                37588 Console                2     243,916 K
python.exe                21464 Console                2      16,760 K
===active temp work dirs (evidence of progress)===
C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/tmp87u1n3pb
===clips in newest tmp===
total 75757
drwxr-xr-x 1 avidu 197609      0 Jul  8 11:47 .
drwxr-xr-x 1 avidu 197609      0 Jul  8 11:47 ..
-rw-r--r-- 1 avidu 197609 39363597 Jul  8 11:47 b_0.mp4
-rw-r--r-- 1 avidu 197609 38199704 Jul  8 11:47 b_1.mp4
-rw-r--r-- 1 avidu 197609      0 Jul  8 11:47 b_2.mp4
-rw-r--r-- 1 avidu 197609     27 Jul  8 11:47 watermark.txt
```

80. assistant (2026-07-08T06:18:31.658Z)

The benchmark is actively progressing ? libx264 4K baseline clips are being encoded (`b\_0`/`b\_1` at ~38 MB each of detailed content, realistic load). Since Python buffers output until exit, I'll get the full results table on completion. This is a good checkpoint to summarize where things stand.

```
**Status so far:**
- **Base reconciled.** #521 (PR #525) is confirmed merged to `origin/master` (`dac6736`). I set up an isolated git **worktree** (`bhi-wt-cheap-watermark`) off master because the shared checkout has a *sibling session's #523 WIP* (including `compiler.py`) that I must not clobber.
- **Found a latent #521 bug** and flagged it (out of scope): `_probe_encoder_available` probes
```

NVENC at `64x64`, below NVENC's minimum encode size, so it **under-reports** NVENC (fails at 64/128px, passes at ?256px) ? silently falling back to slow libx264 on real GPU boxes, possibly even the L4. Chip filed with the concrete fix.

- ***Built an empirical benchmark*** on a real 4K 10-bit HEVC clip (mimicking the corpus decode cost), comparing Baseline / Option A (pre-rendered overlay) / Option B (stream-copy trims + concat re-encode) / Floor (no watermark), across ***both libx264 and h264_nvenc*** (NVENC works at 4K here, so the L4 path is representative). Running now.
- ***Key theoretical finding driving the analysis:*** the watermark burns into **every** frame, so the encode is an unavoidable floor ? Option B **relocates** the re-encode to the concat pass rather than eliminating it, and loses frame-accuracy (keyframe cuts). The benchmark's Floor-vs-Baseline split will quantify exactly how much is irreducible encode vs per-clip overhead, settling A-vs-B with data.

I'll continue automatically when the benchmark completes.

81. user (2026-07-08T06:26:28.218Z)

=== DONE ===

source: 3840x2160, 48.0s | 8 segments | expected reel ~40.0s

encoders under test: ['libx264', 'h264_nvenc']

[libx264	]	BASELINE	total=	73.10s	(trim=	71.26	concat=	1.84)	reel=	40.0s	519.8MB
[libx264	]	OPTION_A	total=	59.52s	(trim=	57.29	concat=	2.07	pre=0.16)	reel=	40.0s
566.7MB											
[libx264	]	OPTION_B	total=	65.29s	(trim=	1.01	concat=	64.28)	reel=	48.1s	568.6MB
[libx264	]	FLOOR	total=	1.10s	(trim=	0.79	concat=	0.31)	reel=	48.1s	308.0MB
[h264_nvenc	]	BASELINE	total=	34.74s	(trim=	32.70	concat=	2.04)	reel=	40.0s	369.4MB
[h264_nvenc	]	OPTION_A	total=	36.02s	(trim=	34.27	concat=	1.59	pre=0.15)	reel=	40.0s
393.8MB											
[h264_nvenc	]	OPTION_B	total=	28.91s	(trim=	0.98	concat=	27.93)	reel=	48.1s	401.8MB
[h264_nvenc	]	FLOOR	total=	1.21s	(trim=	0.86	concat=	0.35)	reel=	48.1s	308.0MB

==== SUMMARY (measured wall seconds) ====

encoder	strategy	total	trim	concat	reel_s	MB
libx264	BASELINE	73.10	71.26	1.84	40.0	519.8
libx264	OPTION_A	59.52	57.29	2.07	40.0	566.7
libx264	OPTION_B	65.29	1.01	64.28	48.1	568.6
libx264	FLOOR	1.10	0.79	0.31	48.1	308.0
h264_nvenc	BASELINE	34.74	32.70	2.04	40.0	369.4
h264_nvenc	OPTION_A	36.02	34.27	1.59	40.0	393.8
h264_nvenc	OPTION_B	28.91	0.98	27.93	48.1	401.8
h264_nvenc	FLOOR	1.21	0.86	0.35	48.1	308.0

82. assistant (2026-07-08T06:28:56.769Z)

Benchmark complete ? clean, decisive data. Let me capture the key findings:

encoder	strategy	total	trim	concat	reel_s	
---	---	---	---	---	---	
libx264	BASELINE	73.1s	71.3	1.8	40.0 ?	
libx264	OPTION_A	**59.5s**	(?18.6%)	57.3	2.1	40.0 ?
libx264	OPTION_B	65.3s	(?10.7%)	1.0	64.3	**48.1 ?**
libx264	FLOOR	1.1s	0.8	0.3	48.1	
nvenc	BASELINE	34.7s	32.7	2.0	40.0 ?	
nvenc	OPTION_A	36.0s	(**+3.7%**)	34.3	1.6	40.0 ?
nvenc	OPTION_B	**28.9s**	(?16.8%)	1.0	27.9	**48.1 ?**
nvenc	FLOOR	1.2s	0.9	0.4	48.1	

Reading it: FLOOR?1.1s proves the stitch is ~99% watermark **encode** (the floor is real). NVENC is the 2.1× lever (#521). ***Option B is frame-INACCURATE*** ? reel is 48.1s vs 40.0s

(stream-copy keyframe cuts add ~8s of unwanted lead-in); that's a content/quality regression, not a cosmetic parity relaxation. Option A is frame-accurate and ?18.6% on libx264 but *+3.7% on NVENC*.

Before deciding, there's a **fourth approach worth measuring**: a **single-pass concat-filter** (multiple `-ss` inputs ? `concat` filter ? watermark ? one encode) ? potentially frame-accurate AND faster on *both* encoders (no intermediate I/O, watermark applied once). If it wins, it beats both A and B. Let me measure it.`

83. assistant (2026-07-08T06:29:23.384Z)

""Focused bench: OPTION_C single-pass concat-filter (frame-accurate, watermark once) vs BASELINE.

OPTION_C: one ffmpeg invocation, N precise `-ss/-t` inputs -> concat filter -> drawtext once -> encode once.`

No intermediate clip files. Should be frame-accurate (reel == 40.0s) unlike stream-copy Option B.

""

```
import os
import subprocess
import sys
import time
import tempfile
```

```
SP = os.path.dirname(os.path.abspath(__file__))
SRC = sys.argv[1] if len(sys.argv) > 1 else os.path.join(SP, "src_4k_hevc10.mp4")
FONT = r"C:/Users/avidu/Projects/khelsutra-guru/bhi-wt-cheap-
watermark/backend/assets/fonts/RobotoSlab-Regular.ttf"
BEGIN_PAD, END_PAD = 0.5, 0.5
RAW_SEGMENTS = [
    (2.0, 6.5), (9.0, 13.5), (15.0, 19.0), (21.5, 25.0),
    (27.0, 31.5), (33.0, 37.0), (39.0, 43.0), (44.0, 47.0),
]
```

```
def ff_quote(p): return p.replace("\\", "/").replace("'", r"\'").replace(":", r"\:")
```

```
def dur(path):
    o = subprocess.run(["ffprobe", "-v", "error", "-show_entries", "format=duration",
                        "-of", "default=nk=1:nw=1", path], capture_output=True, text=True)
    try: return float(o.stdout.strip())
    except Exception: return 0.0
```

```
def drawtext(tmp):
    tf = os.path.join(tmp, "wm.txt")
    open(tf, "w", encoding="utf-8").write("Generated by Khelsutra.guru")
    return (f"drawtext=textfile='{ff_quote(tf)}':fontfile='{ff_quote(FONT)}'"
            f":fontcolor=white@0.85:fontsize=h/29:x=w-tw-24:y=h-th-24"
            f":box=1:boxcolor=black@0.50:boxborderw=10")
```

```
def enc_args(enc):
    return (["-c:v", "libx264", "-preset", "superfast", "-crf", "22"] if enc == "libx264"
            else ["-c:v", enc, "-preset", "p4", "-cq", "23", "-pix_fmt", "yuv420p"])
```

```
def run(cmd):
    t = time.perf_counter()
    p = subprocess.run(cmd, capture_output=True)
    dt = time.perf_counter() - t
    if p.returncode != 0:
        raise RuntimeError(p.stderr.decode(errors="replace")[-1200:])
```

```

return dt

def padded(seg):
    s, e = seg
    return max(0.0, s - BEGIN_PAD), e + END_PAD

def baseline(tmp, enc):
    va, dt = enc_args(enc), drawtext(tmp)
    t0 = time.perf_counter(); clips = []
    for i, seg in enumerate(RAW_SEGMENTS):
        ps, pe = padded(seg); out = os.path.join(tmp, f"b_{i}.mp4")
        run(["ffmpeg", "-y", "-ss", str(ps), "-to", str(pe), "-i", SRC, "-vf", dt,
            *va, "-c:a", "aac", "-b:a", "128k", out]); clips.append(out)
    lp = os.path.join(tmp, "l.txt")
    with open(lp, "w", encoding="utf-8") as f:
        for c in clips: f.write(f"file '%s'\n" % c.replace("\\", "/").replace("'", "'\\'"))
    outp = os.path.join(tmp, "reel_base.mp4")
    run(["ffmpeg", "-y", "-f", "concat", "-safe", "0", "-i", lp, "-c", "copy", outp])
    return time.perf_counter() - t0, outp

def single_pass(tmp, enc):
    """OPTION_C: N -ss/-t inputs -> concat filter -> drawtext once -> one encode."""
    va, dt = enc_args(enc), drawtext(tmp)
    cmd = ["ffmpeg", "-y"]
    for seg in RAW_SEGMENTS:
        ps, pe = padded(seg)
        cmd += ["-ss", str(ps), "-t", str(pe - ps), "-i", SRC]
    n = len(RAW_SEGMENTS)
    fc = []
    for i in range(n):
        fc.append(f"[{i}:v]setpts=PTS-STARTPTS[{i}];[{i}:a]asetpts=PTS-STARTPTS[{i}]")
    concat_in = "".join(f"[v{i}][a{i}]" for i in range(n))
    fc.append(f"{concat_in}concat=n={n}:v=1:a=1[ca]")
    fc.append(f"[cv]{dt}[outv]")
    outp = os.path.join(tmp, "reel_c.mp4")
    t0 = time.perf_counter()
    run(cmd + ["-filter_complex", ";".join(fc), "-map", "[outv]", "-map", "[ca]",
        *va, "-c:a", "aac", "-b:a", "128k", outp])
    return time.perf_counter() - t0, outp

def main():
    encoders = ["libx264"]
    if subprocess.run(["ffmpeg", "-hide_banner", "-f", "lavfi", "-i",
        "testsrc2=s=1920x1080:d=0.2",
        "-frames:v", "3", "-c:v", "h264_nvenc", "-f", "null", "-"],
        capture_output=True).returncode == 0:
        encoders.append("h264_nvenc")
    exp = sum((e + END_PAD) - max(0.0, s - BEGIN_PAD) for s, e in RAW_SEGMENTS)
    print(f"# expected frame-accurate reel ~{exp:.1f}s | encoders {encoders}\n")
    for enc in encoders:
        for name, fn in [("BASELINE", baseline), ("OPTION_C_singlepass", single_pass)]:
            for measured in (False, True):
                with tempfile.TemporaryDirectory(dir=SP) as tmp:
                    try:
                        wall, outp = fn(tmp, enc)
                        d = dur(outp); sz = os.path.getsize(outp) / 1e6
                    except Exception as ex:
                        if measured: print(f"[{enc:11s}] {name:20s} ERROR: {str(ex)[:200]}")
                        break
                    if measured:
                        acc = "ACCURATE" if abs(d - exp) < 0.6 else f"INACCURATE(+{d-exp:.1f}s)"

```

```

print(f"[{enc:11s}] {name:20s} total={wall:7.2f}s reel={d:5.1f}s [{acc}]
{sz:6.1f}MB")

if __name__ == "__main__":
    main()

```

84. user (2026-07-08T06:29:23.447Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\scratchpad\wm_bench2.py (file state is current in your context ? no need to Read it back)

85. assistant (2026-07-08T06:29:30.637Z)

Let me run this focused benchmark (unbuffered so I can watch progress live).

86. user (2026-07-08T06:29:32.888Z)

Command running in background with ID: bzuf204va. Output is being written to: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\tasks\bzuf204va.output. You will be notified when it completes. To check interim output, use Read on that file path.

87. assistant (2026-07-08T06:29:43.721Z)

Let me wait for the single-pass results (first line appears after ~2.5 min ? baseline libx264 warm+measured).

88. user (2026-07-08T06:36:49.572Z)

```

/usr/bin/bash: line 1: [: 0
0: integer expression expected
/usr/bin/bash: line 1: [: 0
0: integer expression expected
/usr/bin/bash: line 1: [: 0
0: integer expression expected
/usr/bin/bash: line 1: [: 0
0: integer expression expected
/usr/bin/bash: line 1: [: 0
0: integer expression expected
/usr/bin/bash: line 1: [: 0
0: integer expression expected
/usr/bin/bash: line 1: [: 0
0: integer expression expected
/usr/bin/bash: line 1: [: 0
0: integer expression expected
/usr/bin/bash: line 1: [: 0
0: integer expression expected
=== COMPLETE ===

```

expected frame-accurate reel ~40.0s | encoders ['libx264', 'h264_nvenc']

[libx264]	BASELINE	total= 78.01s reel= 40.0s [ACCURATE]	519.8MB
[libx264]	OPTION_C_singlepass	total= 67.43s reel= 40.0s [ACCURATE]	490.8MB
[h264_nvenc]	BASELINE	total= 38.49s reel= 40.0s [ACCURATE]	369.4MB
[h264_nvenc]	OPTION_C_singlepass	total= 34.07s reel= 40.0s [ACCURATE]	344.8MB

89. assistant (2026-07-08T06:40:19.043Z)

****OPTION_C (single-pass concat-filter) is the winner**** ? and it beats both A and B:


```

343         f"This segment cannot be extracted ? the video ran out. Skipping."
344     )
345     return None, (idx, start, end, "start beyond video end")
346
347     if end > video_duration:
348         original_end = end
349         end = video_duration
350         logger.warning(
351             f"{segment_label}: end_time ({original_end:.2f}s) exceeds video
duration ({video_duration:.2f}s). "
352             f"Clamping end_time to {video_duration:.2f}s. The segment may be
truncated."
353         )
354
355     clip_path = os.path.join(tmp_dir, f"clip_{idx}.mp4")
356
357     # Precise sub-second trim using fast re-encoding to preserve stream headers
358     padded_start = max(0.0, start - self.begin_padding)
359     padded_end = end + self.end_padding
360
361     # Clamp padded_end to video duration if known
362     if video_duration > 0 and padded_end > video_duration:
363         padded_end = video_duration
364         logger.info(
365             f"{segment_label}: Padded end ({end + self.end_padding:.2f}s) exceeds
video duration. "
366             f"Clamping to {video_duration:.2f}s (end padding partially applied)."
367         )
368
369     trim_duration = padded_end - padded_start
370     logger.info(
371         f"Trimming {segment_label} -> [{padded_start:.3f}s - {padded_end:.3f}s]
(duration: {trim_duration:.2f}s)"
372     )
373
374     # Default to libx264 (the historical settings) when a caller invokes
_trim_one_segment
375     # directly without resolving an encoder ? keeps this method's old behaviour
intact.
376     codec_args = (
377         video_codec_args
378         if video_codec_args is not None
379         else ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"]
380     )
381     base_cmd = [
382         ffmpeg_bin(),
383         "-y",
384         "-ss",
385         str(round(padded_start, 3)),
386         "-to",
387         str(round(padded_end, 3)),
388         "-i",
389         raw_video_path,
390         *codec_args,
391         "-c:a",
392         "aac",
393         "-b:a",
394         "128k",
395     ]
396     vf_args = ["-vf", watermark_filter] if watermark_filter else []
397
398     # Try with the watermark first; if that encode fails (e.g. a bad font/logo
399     # path), retry once WITHOUT it so a branding misconfig can never nuke the

```


I have the exact code. Let me implement now, starting with the config knob, then the compiler refactor + single-pass path.

****1. Config knob**** (`StitchingConfig.single\_pass\_stitch`, default off = parity-safe):

92. user (2026-07-08T06:40:40.051Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\backend\config\models.py has been updated successfully. (file state is current in your context ? no need to Read it back)

93. assistant (2026-07-08T06:40:56.343Z)

****2. Compiler `\_\_init\_\_`**** ? add the `single\_pass` param:

94. user (2026-07-08T06:41:09.246Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

95. assistant (2026-07-08T06:41:46.860Z)

****3. Extract `\_resolve\_padded\_window`**** (shared validation/padding) and refactor `\_trim\_one\_segment` to call it ? a verbatim move, so the per-clip path's behavior is unchanged:

96. user (2026-07-08T06:41:59.811Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

97. assistant (2026-07-08T06:42:06.535Z)

****4. Insert `\_resolve\_padded\_window`**** before `\_trim\_one\_segment`:

98. user (2026-07-08T06:42:20.148Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

99. assistant (2026-07-08T06:43:42.177Z)

****5. Add `\_probe\_has\_audio`**** (ffprobe sibling, needed by the single-pass concat graph):

100. user (2026-07-08T06:43:47.723Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

101. assistant (2026-07-08T06:43:50.809Z)

****6. Add `\_compile\_single\_pass`**** (the core single-pass stitch) before `compile\_highlights`:

102. user (2026-07-08T06:44:17.414Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

103. assistant (2026-07-08T06:44:25.044Z)

****7.** Wire the single-pass branch into `compile\_highlights` (opt-in, with fall-through to the per-clip path):

104. user (2026-07-08T06:44:30.582Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

105. assistant (2026-07-08T06:44:37.825Z)

****8.** Wire config ? `VideoCompiler` in `orchestration.\_stitch\_reel` (rides the cloud compile-spec automatically since it's on `stitching`):

106. user (2026-07-08T06:44:40.727Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

107. assistant (2026-07-08T06:44:45.995Z)

Now let me verify the refactor didn't break the default path. First a syntax check, then the targeted compiler/watermark tests.

108. user (2026-07-08T06:44:50.959Z)

```
SYNTAX OK
===targeted tests===
..... [100%]
30 passed in 0.67s
```

109. assistant (2026-07-08T06:45:13.843Z)

All 30 existing tests pass ? the refactor and the default-off single-pass path preserve the per-clip behavior. Now the critical step: ****verify the single-pass path end-to-end with real ffmpeg**** on the actual 4K source (mocked tests can't prove the `filter\_complex` produces a valid, frame-accurate, watermarked reel). Let me drive the real `VideoCompiler`.

110. assistant (2026-07-08T06:45:44.106Z)

"""End-to-end verification of the single-pass stitch against REAL ffmpeg on the 4K source.

Drives the actual VideoCompiler (single_pass=True) and confirms:

- frame-accurate reel duration (matches per-clip path, within a frame)
- valid 4K video + audio streams
- watermark actually rendered (bottom-right corner differs from a no-watermark reel)
- NVENC single-pass path also works
- fallback to per-clip when the single pass can't run

"""

```
import os
import subprocess
import sys
```

```
WT = r"C:/Users/avidu/Projects/khelsutra-guru/bhi-wt-cheap-watermark"
sys.path.insert(0, WT)
SP = os.path.dirname(os.path.abspath(__file__))
SRC = os.path.join(SP, "src_4k_hevc10.mp4")
OUT = os.path.join(SP, "verify_out")
os.makedirs(OUT, exist_ok=True)
```

```
from backend.pipeline.stitcher.compiler import VideoCompiler # noqa: E402
```

```
SEGMENTS = [(3.0, 7.0), (12.0, 16.0), (20.0, 24.0)]
WM = {"enabled": True, "text": "Generated by Khelsutra.guru"}
```

```

def dur(p):
    o = subprocess.run(["ffprobe", "-v", "error", "-show_entries", "format=duration",
                        "-of", "default=nk=1:nw=1", p], capture_output=True, text=True)
    return float(o.stdout.strip())

def streams(p):
    o = subprocess.run(["ffprobe", "-v", "error", "-show_entries",
                        "stream=codec_type,codec_name,width,height", "-of", "csv=p=0", p],
                        capture_output=True, text=True)
    return o.stdout.strip()

def corner_frame(reel, t, name):
    """Extract the bottom-right 700x180 corner at time t as a PNG (where the watermark sits)."""
    fp = os.path.join(OUT, name)
    subprocess.run(["ffmpeg", "-y", "-ss", str(t), "-i", reel, "-frames:v", "1",
                    "-vf", "crop=700:180:in_w-700:in_h-180", fp], capture_output=True)
    return fp

def mean_luma(png):
    o = subprocess.run(["ffmpeg", "-i", png, "-vf", "format=gray,signalstats", "-f", "null",
                        "-"],
                        capture_output=True, text=True)
    for line in o.stderr.splitlines():
        if "YAVG" in line:
            for tok in line.split():
                if tok.startswith("YAVG:"):
                    return float(tok.split(":")[1])
    return -1.0

exp = sum((e + 1.0) - max(0.0, s - 0.5) for s, e in SEGMENTS) # VideoCompiler defaults 0.5/1.0
print(f"# expected frame-accurate reel = {exp:.2f}s\n")

```

1. per-clip (baseline) WITH watermark ? reference

```

pc = VideoCompiler(OUT, watermark=WM, single_pass=False)
r_pc = pc.compile_highlights(SRC, SEGMENTS, "reel_pc.mp4")
print(f"per-clip      : dur={dur(r_pc):.2f}s streams=[{streams(r_pc)}]")

```

2. single-pass WITH watermark (libx264)

```

spc = VideoCompiler(OUT, watermark=WM, single_pass=True)
r_sp = spc.compile_highlights(SRC, SEGMENTS, "reel_sp.mp4")
print(f"single-pass   : dur={dur(r_sp):.2f}s streams=[{streams(r_sp)}]")

```

3. single-pass WITHOUT watermark

```

spn = VideoCompiler(OUT, watermark={"enabled": False}, single_pass=True)
r_spn = spn.compile_highlights(SRC, SEGMENTS, "reel_sp_nowm.mp4")
print(f"single nowm   : dur={dur(r_spn):.2f}s streams=[{streams(r_spn)}]")

```

4. single-pass NVENC

```

try:
    spv = VideoCompiler(OUT, watermark=WM, single_pass=True, encoder="h264_nvenc")
    r_spv = spv.compile_highlights(SRC, SEGMENTS, "reel_sp_nvenc.mp4")
    print(f"single nvenc : dur={dur(r_spv):.2f}s streams=[{streams(r_spv)}]")
except Exception as ex:
    r_spv = None
    print(f"single nvenc : SKIP/ERROR {str(ex)[:80]}")

```

5. fallback: single_pass=True but a missing source path forces per-clip fallback path.

(Use a valid source but an impossible segment set that yields no windows -> returns falls through to per-clip, which then raises RuntimeError with no clips. We assert the single-pass returned None by observing the per-clip path ran. Simpler: force >cap so

```
spcap = VideoCompiler(OUT, watermark=WM, single_pass=True)
spcap._SINGLE_PASS_MAX_SEGMENTS = 1 # force the cap so single-pass declines -> per-clip
fallback
r_fb = spcap.compile_highlights(SRC, SEGMENTS, "reel_fallback.mp4")
print(f"fallback      : dur={dur(r_fb):.2f}s streams=[{streams(r_fb)}] (single-pass declined,
per-clip ran)")
```

--- assertions ---

```
print("\n# checks:")
ok = True

def check(name, cond):
    global ok
    print(f"  [{ 'PASS' if cond else 'FAIL' }] {name}")
    ok = ok and cond

check(f"single-pass duration ~ per-clip ({dur(r_sp):.2f} vs {dur(r_pc):.2f})", abs(dur(r_sp) -
dur(r_pc)) < 0.6)
check(f"single-pass frame-accurate (~{exp:.1f}s)", abs(dur(r_sp) - exp) < 0.6)
check("single-pass has video+audio", "video" in streams(r_sp) and "audio" in streams(r_sp))
check("single-pass 4K", "3840,2160" in streams(r_sp))
check("no-wm single-pass frame-accurate", abs(dur(r_spn) - exp) < 0.6)
if r_spv:
    check("nvenc single-pass frame-accurate", abs(dur(r_spv) - exp) < 0.6)
check("fallback produced a reel", os.path.exists(r_fb) and abs(dur(r_fb) - exp) < 0.6)
```

watermark visible: bottom-right corner luma differs wm vs no-wm (the white text + box)

```
fp_wm = corner_frame(r_sp, 3.0, "corner_wm.png")
fp_nowm = corner_frame(r_spn, 3.0, "corner_nowm.png")
y_wm, y_nowm = mean_luma(fp_wm), mean_luma(fp_nowm)
check(f"watermark rendered (corner luma wm={y_wm:.1f} vs nowm={y_nowm:.1f}, differ)", abs(y_wm -
y_nowm) > 1.0)

print("\nRESULT:", "ALL PASS" if ok else "SOME FAILED")
print("watermarked-corner PNG for visual check:", fp_wm)
sys.exit(0 if ok else 1)
```

111. user (2026-07-08T06:45:44.145Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\scratchpad\verify_singlepass.py (file state is current in your context ? no need to Read it back)

112. assistant (2026-07-08T06:45:50.167Z)

Let me run the end-to-end verification against real ffmpeg.

113. user (2026-07-08T06:45:52.112Z)

Command running in background with ID: b97k9dhwX. Output is being written to: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\tasks\b97k9dhwX.output. You will be notified when it completes. To check interim output, use Read on that file path.

114. assistant (2026-07-08T06:45:57.783Z)

Let me wait for the verification to complete.

115. user (2026-07-08T06:47:48.993Z)

=== DONE ===

expected frame-accurate reel = 16.50s

```
per-clip      : dur=16.52s streams=[h264,video,3840,2160
aac,audio]
single-pass   : dur=16.50s streams=[h264,video,3840,2160
aac,audio]
single nowm   : dur=16.50s streams=[h264,video,3840,2160
aac,audio]
[compile] encoder 'h264_nvenc' is not runnable in this ffmpeg build ? falling back to libx264
(CPU). On the L4 worker this should be available.
single nvenc  : dur=16.50s streams=[h264,video,3840,2160
aac,audio]
[compile] single-pass stitch unavailable/failed; using the per-clip path.
fallback      : dur=16.52s streams=[h264,video,3840,2160
aac,audio] (single-pass declined, per-clip ran)
```

checks:

```
[PASS] single-pass duration ~ per-clip (16.50 vs 16.52)
[PASS] single-pass frame-accurate (~16.5s)
[PASS] single-pass has video+audio
[PASS] single-pass 4K
[PASS] no-wm single-pass frame-accurate
[PASS] nvenc single-pass frame-accurate
[PASS] fallback produced a reel
[FAIL] watermark rendered (corner luma wm=-1.0 vs nowm=-1.0, differ)
```

RESULT: SOME FAILED

watermarked-corner PNG for visual check: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\scratchpad\verify_out\corner_wm.png

116. assistant (2026-07-08T06:48:07.910Z)

Verification is essentially successful ? ****all substantive checks pass****: single-pass is frame-accurate (16.50s vs per-clip 16.52s), produces valid 4K H.264 + AAC, the no-watermark and NVENC-requested and fallback paths all work. Two notes:

1. The "nvenc" run ****fell back to libx264 via the `\_select\_encoder` probe**** (`64x64` probe bug I flagged) ? the single-pass correctly honored the fallback and still produced a valid reel. My earlier benchmark verified the true NVENC single-pass directly (bypassing the probe): 34.1s, frame-accurate.
2. The only FAIL is my ****luma-check helper**** (`signalstats` YAVG parsing returned ?1.0 for ***both*** ? a broken probe, not a broken watermark).

Let me confirm the watermark actually rendered by visually inspecting the extracted corner frames.

117. assistant (2026-07-08T06:48:59.578Z)

****The watermark renders perfectly**** in the single-pass output ? "...d by Khelsutra.guru" with its backing box, cropped to the bottom-right corner. So the end-to-end verification is a ****success****; the only FAIL was my broken luma-parsing helper, not the watermark.

****Verified with real ffmpeg:**** single-pass is frame-accurate (16.50s vs per-clip 16.52s), valid 4K H.264+AAC, watermark correct, no-watermark path works, per-clip fallback works, and the NVENC request correctly fell back to libx264 via the (#521-buggy) probe while still producing a valid reel.

Now let me add unit tests for the single-pass path (mocked ffmpeg ? command shape, fallback, no-audio, cap) and the extracted validator.

118. user (2026-07-08T06:49:00.843Z)

```
196
197
198     def test_compile_uses_selected_encoder_in_trim(tmp_path, monkeypatch):
199         """End-to-end (ffmpeg mocked): with an available NVENC encoder, every per-clip trim
command
200         carries -c:v h264_nvenc, and the concat stays -c copy."""
201         from backend.pipeline.stitcher import compiler as comp_mod
202
203         monkeypatch.setattr(comp_mod, "_probe_encoder_available", lambda enc: True)
204         comp = VideoCompiler(str(tmp_path / "out"), encoder="h264_nvenc")
205         raw = tmp_path / "raw.mp4"
206         raw.write_bytes(b"x")
207         trims = []
208
209         def fake_run(cmd, *a, **k):
210             if cmd[0] == "ffprobe":
211                 return MagicMock(stdout="100.0\n", returncode=0)
212             out = cmd[-1]
213             if isinstance(out, str) and out.endswith(".mp4"):
214                 open(out, "wb").write(b"clip")
215                 if "-i" in cmd and cmd.index("-i") + 1 == str(raw):
216                     trims.append(cmd)
217             return MagicMock(returncode=0, stdout=b"", stderr=b"")
218
219         with patch.object(subprocess, "run", side_effect=fake_run):
220             comp.compile_highlights(str(raw), [(0.0, 5.0), (10.0, 15.0)], "reel.mp4")
221         assert trims, "no trim commands captured"
222         for cmd in trims:
223             assert "h264_nvenc" in cmd # hardware encode on every clip
224             assert "libx264" not in cmd
225
```

119. assistant (2026-07-08T06:50:04.985Z)

Now let me add unit tests for the single-pass path and the extracted validator, appended to `test\_compiler.py`:

120. user (2026-07-08T06:50:39.921Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\tests\test_compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

121. assistant (2026-07-08T06:50:44.572Z)

Let me run the new tests plus the full compiler/watermark suite:

122. user (2026-07-08T06:50:49.790Z)

```
..... [100%]
37 passed in 0.68s
```

123. assistant (2026-07-08T06:50:57.514Z)

All 37 pass (30 original + 7 new). Let me confirm the config field is tested alongside the existing `encoder` config test, mirroring that pattern:

124. user (2026-07-08T06:50:59.833Z)

```
===where is stitching.encoder config default tested?===
tests/test_compiler.py:258: def test_single_pass_defaults_off(tmp_path):
tests/test_compiler.py:259:     """single_pass defaults OFF ? the per-clip path stays the parity-
safe default."""
```

```

tests/test_compiler.py:260:     assert VideoCompiler(str(tmp_path / "o")).single_pass is False
tests/test_compiler.py:263: def
test_single_pass_fuses_reel_in_one_ffmpeg_pass_with_concat_and_watermark(tmp_path):
tests/test_compiler.py:264:     """single_pass=True: ONE ffmpeg encode via -filter_complex
(concat + drawtext), N precise
tests/test_compiler.py:269:         single_pass=True,
tests/test_compiler.py:305: def
test_single_pass_without_watermark_maps_concat_output_directly(tmp_path):
tests/test_compiler.py:307:         str(tmp_path / "out"), watermark={"enabled": False},
single_pass=True
tests/test_compiler.py:332: def test_single_pass_no_audio_builds_video_only_graph(tmp_path):
tests/test_compiler.py:336:         single_pass=True,
tests/test_compiler.py:361: def test_single_pass_falls_back_to_per_clip_on_failure(tmp_path):
tests/test_compiler.py:367:         single_pass=True,
tests/test_compiler.py:394: def test_single_pass_over_cap_falls_back_to_per_clip(tmp_path):
tests/test_compiler.py:399:         single_pass=True,
tests/test_person_detector.py:118: def test_detections_over_windows_single_pass(mock_cap):
tests/test_worker.py:1004:     seen["encoder"] = stitching.encoder
Binary file tests/__pycache__/test_compiler.cpython-313-pytest-9.0.3.pyc matches

```

125. assistant (2026-07-08T06:51:06.554Z)

Let me check how the worker/dispatch test verifies `encoder` flows through ? since `single\_pass\_stitch` should ride the same `stitching` spec to the L4 worker, I want to confirm that path is covered.

126. user (2026-07-08T06:51:07.328Z)

```

985         json.dump(spec, f)
986     return spec_uri
987
988
989     def test_cmd_compile_stitches_and_writes_marker(tmp_path, monkeypatch):
990         """cmd_compile reads the staged spec, runs the SHARED _stitch_reel core
991         (VideoCompiler stubbed so
992         no ffmpeg), mirrors the reel, and writes a compile done-marker for the dispatcher's
993         bucket-poll."""
994         import backend.orchestration as orch
995
996         bucket = tmp_path / "bucket"
997         bucket.mkdir()
998         source = tmp_path / "src.mp4"
999         source.write_bytes(b"SRC")
1000         spec_uri = _write_compile_spec(bucket, "vidReel", source, "reel.mp4")
1001         done = str(bucket / "_compile" / "tok.done")
1002
1003         seen = {}
1004
1005         def _fake_stitch(cfg, stitching, video_id, source_ref, time_pairs, out_name, locale,
1006         notify):
1007             seen["encoder"] = stitching.encoder
1008             seen["pairs"] = time_pairs
1009             seen["source"] = source_ref
1010             notify("stitching")
1011             return {
1012                 "status": "success",
1013                 "video_id": video_id,
1014                 "download_url": f"/api/highlights/{video_id}/{out_name}",
1015                 "reel_uri": f"gs://b/highlights/{video_id}/{out_name}",
1016             }
1017
1018         monkeypatch.setattr(orch, "_stitch_reel", _fake_stitch)
1019
1020         cfg = tmp_path / "c.json"
1021         cfg.write_text('{"storage": {"backend": "local"}}')

```

```

1019         rc = wmain.cmd_compile(
1020             wmain.build_parser().parse_args(
1021                 [
1022                     "compile",
1023                     "--spec-uri",
1024                     spec_uri,
1025                     "--out-uri",
1026                     str(bucket),
1027                     "--done-uri",
1028                     done,
1029                     "--config",

```

127. assistant (2026-07-08T06:51:20.940Z)

Since `single\_pass\_stitch` rides the compile-spec via `stitching.model\_dump()` (same mechanism as `encoder`, already tested end-to-end), let me confirm that serialization path and see whether a focused spec round-trip assertion is cheap to add.

128. user (2026-07-08T06:51:24.059Z)

```

===spec stitching serialization in orchestration===
103:         config.model_dump(mode="json")
104:         if hasattr(config, "model_dump")
485:             "results": [r.model_dump() for r in val_results],
592:             storage_cfg.model_dump() if storage_cfg is not None else None
613:             video_id, staged_uri, "processed", [r.model_dump() for r in val_results]
702:         storage_cfg.model_dump() # type: ignore[union-attr] # hasattr-guarded; mypy can't
narrow
703:         if hasattr(storage_cfg, "model_dump")
743:             "results": [r.model_dump() for r in val_results],
919:             [r.model_dump() for r in val_results],
929:             [r.model_dump() for r in val_results],

===test_compile_dispatch stitching/spec assertions===
42:         vid = self._video_id or self.spec["video_id"]
43:         out_name = self.spec["output_name"]
125:     and the L4 encoder (deployment.cloud_stitch_encoder = h264_nvenc by default) ? so the
worker
134:     assert fake.spec["video_id"] == video_id
135:     assert fake.spec["video_uri"].endswith("source.mp4")
137:     assert fake.spec["time_pairs"] == [[2.0, 6.0], [10.0, 14.0], [20.0, 22.0]]
138:     # the L4 stitch uses NVENC (forwarded from deployment.cloud_stitch_encoder)
139:     assert fake.spec["stitching"]["encoder"] == "h264_nvenc"
183:     def _fake_stitch(c, stitching, vid, source, pairs, name, locale, notify):

```

129. assistant (2026-07-08T06:51:29.750Z)

Let me see the exact spec-building in orchestration and the dispatch test around line 125-145, to add a `single\_pass\_stitch` round-trip assertion where `encoder` is already checked.

130. user (2026-07-08T06:51:32.629Z)

```

===orchestration: where spec['stitching'] is built===
1122:     stitching = getattr(cfg, "stitching", None) or StitchingConfig()
1135:         "[compile] stitching.min_shot_count=%s: dropped %d/%d segments below the
floor.",
1144:         f"stitching.min_shot_count={stitching.min_shot_count}.",
1203:     with ``stitching.encoder`` (NVENC on the L4, libx264 on the CP), fetches the source,
runs the
1214:     localized_watermark = localize_watermark(stitching.watermark, locale)
1217:     begin_padding=stitching.begin_padding,
1218:     end_padding=stitching.end_padding,
1220:     encoder=getattr(stitching, "encoder", "libx264"),
1221:     single_pass=getattr(stitching, "single_pass_stitch", False),

```



```

1239:         "[compile] stitching failed for video_id=%s: %s", video_id, e, exc_info=True

        video, kept, out_name = _resolve_compile(
            db, cfg, video_id, segment_ids, output_filename
        )
    except _CompileError as e:
        return {"status": "failed", "message": e.detail, "video_id": video_id}

    stitching = getattr(cfg, "stitching", None) or StitchingConfig()
    storage_cfg = getattr(cfg, "storage", None)
    out_root = (getattr(storage_cfg, "output_root", "") or "").strip()
    compute = get_compute_backend(cfg, target_override="cloud_run") if out_root else None
    if compute is None:
        # Compile was placed on the L4 but the cloud dispatch isn't actually wired here (no
output
        # bucket / Cloud Run env). Fall back to the in-process stitch rather than fail ? a slow
reel
        # beats no reel. (The placement guard means this is only reachable in a mis-provisioned
setup.)
        logger.warning(
            "[compile] compile placed on the L4 but no cloud backend available (out_root=%r) ? "
            "stitching in-process on the control-plane instead.",
            out_root,
        )
        time_pairs = [(s["start_time"], s["end_time"]) for s in kept]
        return _stitch_reel(
            cfg, stitching, video_id, video["filepath"], time_pairs, out_name, locale, notify
        )

    # Build the compile-spec the worker consumes. time-pairs are JSON lists (frozen-safe on the
wire);
    # the stitching config is the control-plane's RESOLVED one, with the encoder swapped to the
L4's
    # (deployment.cloud_stitch_encoder, NVENC by default) so the worker hardware-encodes the
reel.
    stitch_dump = stitching.model_dump()
    stitch_dump["encoder"] = getattr(
        getattr(cfg, "deployment", None), "cloud_stitch_encoder", "h264_nvenc"
    )
    spec_payload = {
        "video_id": video_id,
        "video_uri": video["filepath"],
        "output_name": out_name,
        "locale": locale,
        "time_pairs": [
            [float(s["start_time"]), float(s["end_time"])] for s in kept
        ],
        "stitching": stitch_dump,
    }
    token = uuid.uuid4().hex
    spec_uri = f"{out_root.rstrip('/')}/_compile/{token}/spec.json"
    try:
        notify("staging compile spec")
        with tempfile.TemporaryDirectory(prefix="cr-compile-") as td:

```

131. assistant (2026-07-08T06:51:43.693Z)

`single\_pass\_stitch` rides the spec automatically via `stitching.model\_dump()` (line 1326) ? no extra wiring needed. Let me read the dispatch test to lock this round-trip with an assertion:

132. user (2026-07-08T06:51:45.444Z)

```

100         "video_id": video_id,
101         "kind": "compile",
102         "params": json.dumps(

```

```

103         {"segment_ids": segment_ids, "output_filename": name, "locale": locale}
104     ),
105 }
106
107
108 def test_compile_dispatches_to_l4_and_returns_reel(compile_env):
109     db, cfg, video_id, seg_ids, _bucket, monkeypatch = compile_env
110     fake = _FakeCompileClient()
111     _inject_fake(monkeypatch, fake)
112
113     out = orch._execute_compile(cfg, db, _job(video_id, seg_ids, "myreel.mp4"))
114
115     assert out["status"] == "success"
116     assert out["video_id"] == video_id
117     assert out["download_url"] == f"/api/highlights/{video_id}/myreel.mp4"
118     assert out["compute"]["path"] == "cloud_run"
119     assert out["compute"]["execution"] == "exec-compile-1"
120     assert len(fake.calls) == 1 # exactly one Cloud Run Job execution
121
122
123 def test_compile_dispatch_stages_resolved_spec_with_nvenc(compile_env):
124     """The staged compile-spec carries the RESOLVED time-pairs (padding-agnostic
125 start/end from the DB)
126     and the L4 encoder (deployment.cloud_stitch_encoder = h264_nvenc by default) ? so
127 the worker
128     hardware-encodes and needs no control-plane DB."""
129     db, cfg, video_id, seg_ids, _bucket, monkeypatch = compile_env
130     fake = _FakeCompileClient()
131     _inject_fake(monkeypatch, fake)
132
133     orch._execute_compile(cfg, db, _job(video_id, seg_ids))
134
135     assert fake.spec is not None
136     assert fake.spec["video_id"] == video_id
137     assert fake.spec["video_uri"].endswith("source.mp4")
138     # 3 resolved (start, end) pairs from the DB rows
139     assert fake.spec["time_pairs"] == [[2.0, 6.0], [10.0, 14.0], [20.0, 22.0]]
140     # the L4 stitch uses NVENC (forwarded from deployment.cloud_stitch_encoder)
141     assert fake.spec["stitching"]["encoder"] == "h264_nvenc"
142     # the argv is a `compile` subcommand carrying the spec + done marker
143     argv = fake.calls[0]
144     assert argv[0] == "compile"
145     assert "--spec-uri" in argv and "--done-uri" in argv
146     joined = " ".join(argv)

```

133. assistant (2026-07-08T06:51:54.113Z)

Let me add a concise assertion locking that `single\_pass\_stitch` rides the compile-spec to the worker:

134. user (2026-07-08T06:51:57.752Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\tests\test_compile_dispatch.py has been updated successfully. (file state is current in your context ? no need to Read it back)

135. assistant (2026-07-08T06:52:01.141Z)

Now let me run the dispatch/worker tests, then the full suite + lint + types per my verification discipline.

136. user (2026-07-08T06:52:09.706Z)

..... [90%]

[100%]

.....
80 passed in 4.77s

137. assistant (2026-07-08T06:52:16.756Z)

80 pass. Now the full pre-LGTM gate: lint, types, and the whole suite with coverage. Let me check the CI config for the coverage floor and run ruff + mypy on the changed files.

138. user (2026-07-08T06:52:20.164Z)

```
===cov floor / CI config===
.github/workflows/ci.yml:146:      run: python run_all_tests.py --cov=backend --cov=training
--cov-report=term --cov-fail-under=70
.github/workflows/nightly.yml:141:    run: python run_all_tests.py --cov=backend
--cov=training --cov-report=term --cov-fail-under=70

===ruff===
All checks passed!
exit: 0
```

139. assistant (2026-07-08T06:52:23.843Z)

Ruff clean; coverage floor is 70. Let me run mypy (matching CI) and check its config:

140. user (2026-07-08T06:53:08.949Z)

```
===mypy invocation in CI===
.github/workflows/ci.yml:66:  # Type gate (mypy.ini at repo root): lenient green baseline with
an explicit
.github/workflows/ci.yml:68:  # up. Typed core deps are installed so mypy sees their real
signatures
.github/workflows/ci.yml:79:    - name: Install mypy + typed core deps
.github/workflows/ci.yml:80:    # mypy pinned alongside ruff so a type-checker release can't
redde CI overnight.
.github/workflows/ci.yml:81:    run: python -m pip install mypy==2.1.0 pydantic fastapi
uvicorn python-dotenv google-genai numpy
.github/workflows/ci.yml:83:    run: mypy backend training
.github/workflows/ci.yml:136:   # (pytest, pytest-cov, httpx for the fastapi TestClient,
ruff, mypy).
.github/workflows/nightly.yml:74: # Type gate (mypy.ini at repo root) ? mirrors
ci.yml::typecheck. Typed core
.github/workflows/nightly.yml:75: # deps are installed so mypy sees real signatures (without
them
.github/workflows/nightly.yml:86:    - name: Install mypy + typed core deps

===run mypy on changed files===
backend\annotations\base.py:88: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\sports\__init__.py:9: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\pipeline\validators\base.py:123: note: By default the bodies of untyped functions are
not checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 3 source files
```

141. assistant (2026-07-08T06:53:22.733Z)

mypy clean on changed files (the notes are unrelated info). Let me run the full CI gates: `mypy backend training`, `ruff check .`, then the whole test suite with coverage.

142. user (2026-07-08T06:53:31.429Z)

```
===full ruff===
```

All checks passed!

```
===full mypy (backend training)===
backend\annotations\base.py:88: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\sports\__init__.py:9: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\pipeline\validators\base.py:123: note: By default the bodies of untyped functions are
not checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 150 source files
```

143. assistant (2026-07-08T06:53:36.761Z)

Full ruff + mypy clean (150 files). Now the full test suite with coverage (matching CI). Running in background as it's the heavy gate.

144. user (2026-07-08T06:53:37.975Z)

Command running in background with ID: bzpqhwdgh. Output is being written to:
C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\tasks\bpqhwdgh.output. You will be notified when it completes. To check interim output, use Read on that file path.

145. assistant (2026-07-08T06:53:55.340Z)

While the full suite runs, let me write the required documentation. Let me re-read the actual S3 in the worktree to edit it precisely.

146. user (2026-07-08T06:53:56.334Z)

```
100    ---
101
102    ## 3. NVENC on the L4 ? and the D5/D16(b) parity note
103
104    The dominant term in the CUJ stitch was the per-clip 4K re-encode the branding
watermark forces
105    (~67 s/clip on 2 vCPU with libx264). The L4 has a hardware H.264/HEVC encoder, so #521
adds an encoder
106    knob:
107
108    - `StitchingConfig.encoder` ? `libx264` (default, CPU, works everywhere incl. CI) |
`h264_nvenc` |
109    `hevc_nvenc`. `VideoCompiler` probes the requested encoder once per run (a 1-frame
lavfi test
110    encode) and falls back to libx264 if this ffmpeg build can't run it, so a
misconfig / non-GPU
111    box never breaks. All clips share the one chosen encoder ? the concat stays -c copy.
112    - `deployment.cloud_stitch_encoder` (default h264_nvenc) ? what the compile-dispatch
forwards onto
113    the worker's stitching.encoder for the L4 run. Set it to libx264 to run the L4
stitch on CPU.
114    - The control-plane's own local stitch always uses stitching.encoder (libx264 by
default) ? the
115    CP has no GPU.
116
117    Parity note (honest): [README.md](./README.md) D5/D16(b) kept stitch encode =
libx264 to
118    protect a byte-exact stitch parity guarantee, and explicitly deferred NVENC encode
behind a config
119    knob, "revisit only if M7 cost data justifies relaxing parity." The 2026-07-08 CUJ
(~13 min stitch,
120    a hard scaling wall past ~5 concurrent users) is that cost data. #521 lands exactly
```

the deferred

```

121     shape ? NVENC **behind a config knob**, libx264 the default, config-reversible per
phase. Enabling
122     NVENC relaxes byte-identity **at the reel ENCODE layer only** (a cosmetic, shareable
artifact); the
123     **WINDOW/rally-level parity** (the analytical output) is untouched ? detection and
windowing are
124     unchanged. If bit-exact reels are ever required, set `cloud_stitch_encoder=libx264`.
125
126     > Follow-up (tracked, not in #521's first cut): a **pre-rendered / overlay-filter
watermark** so the
127     > stitch isn't a full re-encode at all (the deeper win beyond swapping the encoder).
128
129     ---

```

147. assistant (2026-07-08T06:54:10.993Z)

Now let me replace the deferred follow-up note with the completed, measured Part-3 investigation:

148. user (2026-07-08T06:54:45.669Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\docs\COMPUTE_DECOUPLED_SERVING\PHASE_PLACEMENT.md has been updated successfully. (file state is current in your context ? no need to Read it back)

149. assistant (2026-07-08T06:54:53.408Z)

Let me update the \$5 Files & tests table to include the Part-3 surface. Let me read it first.

150. user (2026-07-08T06:54:54.722Z)

```

205     **no** `--skip-validation`, so the **worker** validates and rejects a bad clip via a
non-zero rc
206     (surfaced as a failed job).
207
208     **Footgun guard:** a clip can only be validated on the worker if the worker also
**detects** it there.
209     Two protections:
210
211     1. **Config-load guard** (`DeploymentConfig` model-validator): `validation=cloud_run_l4`
requires the
212     effective segment placement to also be `cloud_run_l4` (explicit, or via
`compute_target=cloud_run`)
213     ? else config load fails with an actionable error.
214     2. **Runtime guard** (`_will_detect_on_cloud`): if a per-request `compute='local'`
forces detection
215     in-process (no worker), the control-plane **still validates** ? never leave a clip
unvalidated
216     because nobody ran the gate.
217
218     ---
219
220     ## 5. Files & tests
221
222     | Area | Code | Tests |
223     |---|---|---|
224     | Placement resolver + config | `backend/config/placement.py`,
`backend/config/models.py` (`PhasePlacementConfig`, `DeploymentConfig.phase_placement` /
`cloud_stitch_encoder`, guard) | `tests/test_phase_placement.py`,
`tests/test_config_deployment.py` |
225     | Preset | `backend/config/presets.py` (`cloud-serving` ?
`phase_placement.compile=cloud_run_l4`) | `tests/test_phase_placement.py` |
226     | compile dispatch | `backend/orchestration.py` (`_execute_compile`,

```

```
`_dispatch_compile_remote`, `_stitch_reel`, `_will_detect_on_cloud`) |  
`tests/test_compile_dispatch.py` |
```

151. user (2026-07-08T06:55:01.714Z)

```
226 | compile dispatch | `backend/orchestration.py` (`_execute_compile`,  
`_dispatch_compile_remote`, `_stitch_reel`, `_will_detect_on_cloud`) |  
`tests/test_compile_dispatch.py` |  
227 | compute backend | `backend/compute/base.py` (`CompileJobSpec`, `run_compile`),  
`backend/compute/cloud_run.py` (`run_compile`, `_await_marker`) |  
`tests/test_compute_backend.py` |  
228 | worker | `backend/worker/__main__.py` (`compile` subcommand, `cmd_compile`) |  
`tests/test_worker.py` |  
229 | NVENC encoder | `backend/pipeline/stitcher/compiler.py` (`encoder`, `_select_encoder`,  
`_probe_encoder_available`) | `tests/test_compiler.py` |  
230 | validation delegation | `backend/orchestration.py` (`_maybe_dispatch_remote` skip  
flag, `_execute_processing`) | `tests/test_api_cloud_dispatch.py` |  
231  
232 All GPU-free + cloud-free: a fake `CloudRunJobClient` simulates the worker (reads the  
staged spec,  
233 writes the reel + done-marker into a local "bucket"), mirroring  
`tests/test_api_cloud_dispatch.py`.  
234
```

152. user (2026-07-08T06:55:09.971Z)

The file C:\Users\avidu\Projects\khelsutra-guru\bhi-wt-cheap-watermark\docs\COMPUTE_DECOUPLED_SERVING\PHASE_PLACEMENT.md has been updated successfully. (file state is current in your context ? no need to Read it back)

153. assistant (2026-07-08T06:55:12.226Z)

Let me check the full test suite result.

154. user (2026-07-08T06:57:55.271Z)

=== output ===

backend\pipeline\validators\base.py	85	3	96%
backend\pipeline\validators\frame_sampler.py	54	1	98%
backend\providers__init__.py	16	1	94%
backend\providers\base.py	9	1	89%
backend\providers\gemini.py	137	21	85%
backend\reporting__init__.py	26	2	92%
backend\reporting\harvest.py	111	8	93%
backend\reporting\run_signals.py	14	0	100%
backend\results.py	27	0	100%
backend\server_runtime.py	57	2	96%
backend\sports__init__.py	17	1	94%
backend\sports\badminton.py	13	1	92%
backend\sports\base.py	12	3	75%
backend\storage__init__.py	85	3	96%
backend\storage\base.py	16	1	94%
backend\storage\capabilities.py	59	10	83%
backend\storage\gcs.py	80	3	96%
backend\storage\gdrive.py	61	3	95%
backend\storage\gdrive_api.py	105	6	94%
backend\storage\local.py	49	3	94%
backend\storage\ref.py	49	0	100%
backend\storage\s3.py	47	5	89%
backend\tools__init__.py	0	0	100%
backend\tools\export.py	75	2	97%
backend\tools\rally_slicer.py	105	47	55%
backend\utils__init__.py	0	0	100%
backend\utils\annotation_proxy.py	64	11	83%

backend\utils\app_home.py	40	0	100%
backend\utils\ffmpeg.py	47	3	94%
backend\utils\freespace.py	32	2	94%
backend\utils\geometry.py	29	0	100%
backend\utils\hashing.py	9	0	100%
backend\utils\logging_setup.py	11	0	100%
backend\utils\paths.py	44	3	93%
backend\utils\redact.py	16	0	100%
backend\utils\run_telemetry.py	126	12	90%
backend\utils\single_instance.py	43	11	74%
backend\utils\video.py	34	3	91%
backend\utils\workspace.py	49	0	100%
backend\worker__init__.py	0	0	100%
backend\worker__main__.py	371	66	82%
training__init__.py	0	0	100%
training\gen0__init__.py	0	0	100%
training\gen0\harness.py	130	9	93%

TOTAL 12670 1950 85%

Required test coverage of 70% reached. Total coverage: 84.61%

===== short test summary info =====

FAILED tests/test_import_smoke.py::test_backend_package_has_no_syntax_errors
 FAILED tests/test_light_tier_no_torch.py::test_light_tier_imports_without_torch
 FAILED tests/test_main_module_budget.py::test_main_py_line_budget - Assertion...
 3 failed, 1879 passed, 8 skipped, 25 warnings in 250.92s (0:04:10)

=====

RESULT: FAIL - pytest exit 1

=====

155. assistant (2026-07-08T06:58:58.027Z)

?? **Stop ? the worktree state changed underneath me.** The post-run notes show my worktree's `compiler.py` now contains a `\_PROBE\_SOURCE\_SIZE = "256x256"` NVENC-probe fix I did **not** write (that's the separate task I *flagged*), and the shown `\_\_init\_\_` appears to **lack my `single\_pass` param**. Meanwhile `test\_backend\_package\_has\_no\_syntax\_errors` FAILED ? a syntax error somewhere in `backend/`. A concurrent process may have written into my supposedly-isolated worktree. I must assess the actual current state before doing anything else.

156. user (2026-07-08T06:59:04.452Z)

===git status===

===my markers present in compiler.py?===

single_pass-related lines: 0

===external NVENC fix present?===

37:_PROBE_SOURCE_SIZE = "256x256"

42: lavfi test encode to null at ``\_PROBE\_SOURCE\_SIZE`` (a resolution ? NVENC's minimum, so the

62: f"color=c=black:s={_PROBE_SOURCE_SIZE}:d=0.1",

===syntax check whole backend===

no syntax errors in backend/

157. assistant (2026-07-08T06:59:40.532Z)

My uncommitted changes to `compiler.py` are gone ? `git status` is clean, 0 `single\_pass` markers remain, but the external NVENC fix (`\_PROBE\_SOURCE\_SIZE`) is present. A concurrent process wrote into my supposedly-isolated worktree. I need the full picture before recovering ? what happened to HEAD, and whether my other files survived.

158. user (2026-07-08T06:59:45.581Z)

===HEAD + branch===

```
codex/issue-523-sd-reel-review
65ee4c6 fix(serving): probe NVENC at 256x256, not 64x64, so it isn't falsely rejected
(#521/#525) (#527)
008a5d4 fix(compile): fail the L4-worker stitch when the reel mirror-upload fails (#521 follow-
up) (#528)
dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521) (#525)
bf97700 docs(runlogs): record live SPA CUJ closing the 2 owner-gated validation-latency cells
($9) (#522)
```

===reflog (last 8)===

```
65ee4c6 HEAD@{0}: reset: moving to 65ee4c64dcab41c74f63613c973e9069d0167b8d
65ee4c6 HEAD@{1}: checkout: moving from perf/521-cheap-watermark to codex/issue-523-sd-reel-
review
65ee4c6 HEAD@{2}: pull --ff-only: Fast-forward
dac6736 HEAD@{3}: reset: moving to HEAD
dac6736 HEAD@{4}: reset: moving to HEAD
dac6736 HEAD@{5}:
```

===my markers across files===

```
backend/pipeline/stitcher/compiler.py: 0 single_pass hits
backend/config/models.py: 0 single_pass hits
backend/orchestration.py: 0 single_pass hits
tests/test_compiler.py: 0 single_pass hits
tests/test_compile_dispatch.py: 0 single_pass hits
docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md: 0 single_pass hits
```

159. assistant (2026-07-08T07:01:08.857Z)

Now I understand the collision. The reflog shows a **concurrent session hijacked my worktree**:
`checkout: moving from perf/521-cheap-watermark to codex/issue-523-sd-reel-review`, then `reset:
moving to 65ee4c6`. My uncommitted single-pass work (never committed) was wiped from disk.
Notably, **my flagged NVENC-probe bug is now fixed and merged (#527)**, and #528 also landed.

Two lessons: (1) I have every edit in my conversation history, so I can reconstruct; (2) I must
commit immediately after re-applying so a working-tree reset can't destroy it again, and
work in a **scratchpad-isolated** worktree siblings can't discover. Let me re-sync and assess
the new base.

160. user (2026-07-08T07:01:16.228Z)

===origin/master tip===

```
65ee4c6 fix(serving): probe NVENC at 256x256, not 64x64, so it isn't falsely rejected
(#521/#525) (#527)
008a5d4 fix(compile): fail the L4-worker stitch when the reel mirror-upload fails (#521 follow-
up) (#528)
dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521) (#525)
```

===#527/#528 on origin/master?===

```
65ee4c6(#527) IS on origin/master
```

===did #527/#528 touch my files (vs dac6736)?===

```
backend/orchestration.py | 49 ++++++++--
backend/pipeline/stitcher/compiler.py | 19 +++-
docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md | 8 +-
tests/test_compile_dispatch.py | 85 ++++++++
tests/test_compiler.py | 62 ++++++++
5 files changed, 211 insertions(+), 12 deletions(-)
```

===my leftover branch?===

```
perf/521-cheap-watermark
```

161. assistant (2026-07-08T07:02:02.769Z)

Now I have the recovery picture. origin/master advanced to `65ee4c6` (#525 + #528 + **#527 = my flagged NVENC fix, now merged**). My changes were never committed, so they're gone from disk ? but I have every edit in history. #527/#528 touched some of the same files, so I'll re-apply against the **new** base and **commit immediately** this time. I'll work in a **scratchpad-isolated** worktree that sibling automations can't discover.

162. user (2026-07-08T07:02:10.171Z)

Exit code 1

===models.py touched by #527/#528?===

models.py UNCHANGED since dac6736 (safe verbatim re-apply)

===create fresh isolated worktree in scratchpad off origin/master===

error: unable to create file docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-serving_gpu-resident-nvdec_2026-07-07.md: Filename too long

error: unable to create file docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md: Filename too long

fatal: Could not reset index file to revision 'HEAD'.

/usr/bin/bash: line 1: cd: C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer/46eabfd2-0e71-4ff5-a68d-c76f474d4ff3/scratchpad/wt-sp: No such file or directory

163. assistant (2026-07-08T07:02:23.265Z)

Windows MAX_PATH ? the scratchpad path is too long for the repo's deep doc filenames. The existing worktrees used short `C:/kg-wt\*` roots for this reason. Let me clean up the partial worktree and use a short path.

164. user (2026-07-08T07:02:34.178Z)

===prune broken worktree===

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/a1-master-worktree': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/bhi-nvenc-probe-wt': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/bhi-wt-worker-reval': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/fix-worktree': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/kg-wt521': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/kg-wt525fix': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/master-compare': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/pr465': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/pr465-current-merge': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/pr465-merge': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/pr467': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/pr469': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-465fix': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-a1-runlog': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-claudemd-r25': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-cloudrun-cancel': Permission denied

error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-cloudrun-dispatch': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-cp-rebuild-spec': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-dockerfile-auth': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-docs': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-docs-reconcile': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-drain-concurrency': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-ffmpeg-timeouts': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-gdrive-api-safety': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-idempotent-process': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-input-jail': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-output-dir-config': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-p1-genyaml': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-p2-upload-stream': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-p4-md5cache': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-p5a-pollknob': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-p6-heartbeat': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-pr481-review': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-retention-sweep': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-review-492-mergecheck': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-review-493-mergecheck': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-scratch-reclaim': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-signed-download': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-source-cache': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-upload-dir': Permission denied
error: failed to delete 'C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/.git/worktrees/wt-upload-staging': Permission denied
Deleted branch perf/521-single-pass-stitch (was 65ee4c6).

===create at short path C:/kg-wt-sp===
Preparing worktree (new branch 'perf/521-single-pass-stitch')
branch 'perf/521-single-pass-stitch' set up to track 'origin/master'.
HEAD is now at 65ee4c6 fix(serving): probe NVENC at 256x256, not 64x64, so it isn't falsely rejected (#521/#525) (#527)
===verify===
perf/521-single-pass-stitch
65ee4c6 fix(serving): probe NVENC at 256x256, not 64x64, so it isn't falsely rejected (#521/#525) (#527)
(clean if empty)
has #527 probe fix
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

165. assistant (2026-07-08T07:03:22.337Z)

New isolated worktree ready at `C:/kg-wt-sp` on `perf/521-single-pass-stitch` @ 65ee4c6 (latest master with #527/#528), clean tree. (The prune "permission denied" noise is just other sessions' locked worktrees ? this environment has dozens of concurrent worktrees, which explains the hijack.) `models.py` is unchanged since my base; the others were touched by #527/#528, so I'll re-apply against the new content and **commit immediately**.

Let me read the new base regions I need to re-apply into.

166. user (2026-07-08T07:03:22.919Z)

```
1     import logging
2     import os
3     import subprocess
4     import tempfile
5     from pathlib import Path
6     from typing import Any, Dict, List, Optional, Tuple, Union
7
8     from backend.utils.ffmpeg import ( # P2a: single ffmpeg/ffprobe resolver
9         ffmpeg_bin,
10        ffmpeg_timeout_sec,
11        ffprobe_bin,
12        ffprobe_timeout_sec,
13        timeout_label,
14    )
15
16    logger = logging.getLogger(__name__)
17
18    # Bundled, commercially-clean default watermark font (Apache-2.0). Shipped so the
overlay
19    # renders out-of-the-box without a system fontconfig setup (a common gap on Windows,
where
20    # `font='Sans'` fails). Overridable via stitching.watermark.font_path. This module lives
at
21    # backend/pipeline/stitcher/, so parents[2] == backend/.
22    _BUNDLED_FONT = (
23        Path(__file__).resolve().parents[2] / "assets" / "fonts" / "RobotoSlab-Regular.ttf"
24    )
25
26    #: Per-encoder availability memo (#521): probing ffmpeg is a subprocess, so a run's
first clip pays
27    #: it and the rest reuse the verdict. Keyed by encoder name; cleared only across
processes.
28    _ENCODER_AVAILABLE: Dict[str, bool] = {}
29
30    #: Frame geometry for the 1-frame NVENC probe encode in ``_probe_encoder_available``. It
MUST stay
31    #: at or above NVENC's minimum encode dimensions: on real GPUs the hardware H.264/HEVC
encoders
32    #: REJECT tiny frames, so a sub-minimum probe (the original 64x64) FALSELY reports NVENC
as
33    #: unavailable and ``_select_encoder`` silently falls back to libx264 ? defeating the
NVENC-on-L4
34    #: re-encode win of #521/#525. Empirically (RTX 2070 Super, ffmpeg 8.1.1) h264_nvenc
rejects
35    #: 64x64 and 128x128 but accepts ~145px and up; 256x256 keeps a safe margin. Do NOT
shrink below
36    #: 256x256 ? guarded by tests/test_compiler.py::test_probe_encoder_uses_safe_resolution.
37    _PROBE_SOURCE_SIZE = "256x256"
38
39
40    def _probe_encoder_available(encoder: str) -> bool:
41        """Can this ffmpeg build actually RUN ``encoder`` (not merely list it)? Runs a tiny
1-frame
```

```

42     lavfi test encode to null at ``_PROBE_SOURCE_SIZE`` (a resolution ? NVENC's minimum,
so the
43     probe reflects real hardware runnability rather than rejecting an undersized frame).
NVENC is
44     often *compiled in* on boxes without an NVIDIA GPU, so a real test encode ? not
``-encoders``
45     grep ? is what distinguishes "usable" from "will fail at runtime". Memoized. Any
error /
46     timeout ? unavailable (caller falls back to libx264)."""
47     if encoder in _ENCODER_AVAILABLE:
48         return _ENCODER_AVAILABLE[encoder]
49     ok = False
50     # The probe is a tiny 1-frame encode ? bound it to a short timeout (? the configured
ffmpeg
51     # budget, or 30 s when unbounded) so a wedged encoder can't hang the whole stitch.
52     configured = ffmpeg_timeout_sec()
53     probe_timeout = min(30.0, float(configured)) if configured else 30.0
54     try:
55         proc = subprocess.run(
56             [
57                 ffmpeg_bin(),
58                 "-hide_banner",
59                 "-f",
60                 "lavfi",
61                 "-i",
62                 f"color=c=black:s={_PROBE_SOURCE_SIZE}:d=0.1",
63                 "-frames:v",
64                 "1",
65                 "-c:v",
66                 encoder,
67                 "-f",
68                 "null",
69                 "-",
70             ],
71             capture_output=True,
72             timeout=probe_timeout,
73         )
74         ok = proc.returncode == 0
75     except Exception: # noqa: BLE001 ? a probe failure/timeout just means "not usable ?
fall back"
76         ok = False
77     _ENCODER_AVAILABLE[encoder] = ok
78     return ok
79
80
81     class VideoCompiler:
82     def __init__(
83         self,
84         output_dir: str,
85         end_padding: float = 1.0,
86         begin_padding: float = 0.5,
87         watermark: Optional[Any] = None,
88         encoder: str = "libx264",
89     ):
90         self.output_dir = output_dir
91         self.end_padding = end_padding
92         self.begin_padding = begin_padding
93         # Optional branding overlay. Accepts a WatermarkConfig model or a plain dict;
94         # normalized to a dict (or None) so this module stays config-package-agnostic.
95         self.watermark = self._normalize_watermark(watermark)
96         # Video encoder for the per-clip trim re-encode (#521). "libx264" (CPU) is the
default and
97         # works everywhere; "h264_nvenc"/"hevc_nvenc" use the L4's HARDWARE encoder for
the 4K
98         # re-encode the watermark forces. The requested encoder is PROBED once per run

```

```

99         # (_select_encoder) and falls back to libx264 if this ffmpeg build can't run it.
100         self.encoder = (encoder or "libx264").strip() or "libx264"
101         os.makedirs(output_dir, exist_ok=True)
102
103     @staticmethod
104     def _normalize_watermark(watermark: Optional[Any]) -> Optional[Dict[str, Any]]:
105         if watermark is None:
106             return None
107         if hasattr(watermark, "model_dump"):
108             return watermark.model_dump()
109         if isinstance(watermark, dict):
110             return dict(watermark)
111         return None
112
113     @staticmethod
114     def _ff_quote(path: str) -> str:
115         """Make a filesystem path safe to embed inside single quotes in an ffmpeg
116         filtergraph operand (drawtext fontfile/textfile, movie= source): forward
117         slashes (Windows-friendly), escaped single quotes, AND escaped colons.
118
119         The colon MUST be escaped even inside single quotes: ffmpeg's filtergraph
120         option parser treats ':' as the option separator, so a Windows drive-letter
121         path like ``C:/dir/font.ttf`` otherwise truncates the operand at ``C`` and
122         the encode fails with "Invalid argument". This silently disabled the
123         on-by-default bundled-font watermark (and absolute logo paths) on Windows."""
124         return path.replace("\\", "/").replace("'", r"\'").replace(":", r"\:")
125
126     @staticmethod
127     def _parse_time(val: Union[int, float, str]) -> float:
128         """Normalizes a timestamp value to float seconds.
129         Handles: float/int (pass-through), plain numeric strings ("12.4"),
130         and colon-separated strings ("MM:SS", "HH:MM:SS", "H:MM:SS.ms").
131         This mirrors the parse_time logic used in gemini.py to ensure
132         the stitcher can always consume whatever format the indexer produced."""
133         if isinstance(val, (int, float)):
134             return float(val)
135         if isinstance(val, str):
136             val = val.strip()
137             if ":" in val:
138                 # Handle MM:SS, HH:MM:SS, etc.
139                 parts = val.split(":")
140                 parts.reverse()
141                 total = 0.0
142                 for i, p in enumerate(parts):
143                     try:
144                         total += float(p) * (60**i)
145                     except ValueError:
146                         pass
147                 return total
148             try:
149                 return float(val)
150             except ValueError:
151                 logger.warning(
152                     f"Could not parse timestamp value '{val}' as a number. Defaulting to
0.0."
153                 )
154                 return 0.0
155             logger.warning(
156                 f"Unexpected timestamp type {type(val).__name__} for value '{val}'.
Defaulting to 0.0."
157             )
158             return 0.0
159
160     @staticmethod
161     def _merge_overlapping(

```

```

162         segments: List[Tuple[float, float]],
163     ) -> List[Tuple[float, float]]:
164         """Normalize, sort, and merge overlapping/duplicate (start, end) ranges.
165
166         A highlight reel must never replay the same footage. The same rally can land in
167         the segment list more than once ? e.g. two detector runs accumulated for one
168         video (add_video is an UPSERT, so re-running without a prune leaves both sets),
169         or
170         chunk-overlap duplicates near boundaries. Stitching those verbatim plays the
171         rally
172         twice. Merging any range that overlaps or abuts the previous one collapses true
173         duplicates/overlaps while preserving genuinely distinct (gapped) rallies.
174         Degenerate rows (end <= start) are dropped. Timestamps may be float/int/str.
175         """
176         parsed = []
177         for raw_s, raw_e in segments:
178             s = VideoCompiler._parse_time(raw_s)
179             e = VideoCompiler._parse_time(raw_e)
180             if e > s:
181                 parsed.append((s, e))
182         parsed.sort()
183         merged: List[Tuple[float, float]] = []
184         for s, e in parsed:
185             if merged and s <= merged[-1][1]: # overlaps/abuts previous -> extend it
186                 merged[-1] = (merged[-1][0], max(merged[-1][1], e))
187             else:
188                 merged.append((s, e))
189         return merged
190
191 def _get_video_duration(self, video_path: str) -> float:
192     """Probes the video file to get its total duration in seconds.
193     Returns 0.0 if it cannot be determined."""
194     try:
195         timeout = ffprobe_timeout_sec()
196         result = subprocess.run(
197             [
198                 ffprobe_bin(),
199                 "-v",
200                 "error",
201                 "-show_entries",
202                 "format=duration",
203                 "-of",
204                 "default=noprint_wrappers=1:nokey=1",
205                 video_path,
206             ],
207             capture_output=True,
208             text=True,
209             check=True,
210             timeout=timeout,
211         )
212         return float(result.stdout.strip())
213     except subprocess.TimeoutExpired:
214         logger.warning(
215             "Could not determine video duration via ffprobe: timed out after %s",
216             timeout_label(ffprobe_timeout_sec()),
217         )
218         return 0.0
219     except Exception as e:
220         logger.warning(f"Could not determine video duration via ffprobe: {e}")
221         return 0.0
222
223 def _watermark_filter(self, tmp_dir: str) -> Optional[str]:
224     """Build the ffmpeg -vf filtergraph that burns the configured watermark (text
225     and/or logo PNG) into a clip during the per-clip re-encode. Returns None when
226     watermarking is disabled or has nothing to draw. The concat stage stays -c copy

```

```

225         because the mark is already baked into every clip identically.""
226         wm = self.watermark
227         if not wm or not wm.get("enabled"):
228             return None
229
230         text = (wm.get("text") or "").strip()
231         logo_path = (wm.get("logo_path") or "").strip()
232         if logo_path and not os.path.exists(logo_path):
233             logger.warning(
234                 f"[watermark] logo not found: {logo_path} ? skipping logo overlay."
235             )
236             logo_path = ""
237         if not text and not logo_path:
238             return None
239
240         margin = int(wm.get("margin", 24))
241         opacity = float(wm.get("opacity", 0.85))
242         position = str(wm.get("position", "bottom-right"))
243         # drawtext uses lowercase w/h (main) + tw/th (text); overlay uses W/H (main) +
w/h (logo).
244         dt_xy = {
245             "bottom-right": f"x=w-tw-{margin}:y=h-th-{margin}",
246             "bottom-left": f"x={margin}:y=h-th-{margin}",
247             "top-right": f"x=w-tw-{margin}:y={margin}",
248             "top-left": f"x={margin}:y={margin}",
249         }.get(position, f"x=w-tw-{margin}:y=h-th-{margin}")
250         ov_xy = {
251             "bottom-right": f"W-w-{margin}:H-h-{margin}",
252             "bottom-left": f"{margin}:H-h-{margin}",
253             "top-right": f"W-w-{margin}:{margin}",
254             "top-left": f"{margin}:{margin}",
255         }.get(position, f"W-w-{margin}:H-h-{margin}")
256
257         drawtext = ""
258         if text:
259             ratio = float(wm.get("font_size_ratio") or 0.035)
260             divisor = max(1, round(1.0 / ratio))
261             # textfile= sidesteps the notorious drawtext text= escaping for arbitrary
strings.
262             tf_path = os.path.join(tmp_dir, "watermark.txt")
263             with open(tf_path, "w", encoding="utf-8") as fh:
264                 fh.write(text)
265             # Font precedence: explicit font_path ? bundled Apache-2.0 default (renders
266             # without fontconfig) ? fontconfig family (last-resort, environment-
dependent).
267             font_path = (wm.get("font_path") or "").strip()
268             if not font_path and _BUNDLED_FONT.exists():
269                 font_path = str(_BUNDLED_FONT)
270             if font_path:
271                 font_opt = f"fontfile='{self._ff_quote(font_path)}'"
272             else:
273                 font_opt = f"font='{wm.get('font', 'Sans')}'"
274             box = ""
275             if wm.get("box", True):
276                 box = f"box=1:boxcolor=black@{min(opacity, 0.5):.2f}:boxborderw=10"
277             drawtext = (
278                 f"drawtext=textfile='{self._ff_quote(tf_path)}':{font_opt}"
279                 f"fontcolor=white@{opacity:.2f}:fontsize=h/{divisor}:{dt_xy}{box}"
280             )
281
282         if logo_path and drawtext:
283             return
f"movie='{self._ff_quote(logo_path)}'[wm];[in][wm]overlay={ov_xy},{drawtext}"
284         if logo_path:
285             return f"movie='{self._ff_quote(logo_path)}'[wm];[in][wm]overlay={ov_xy}"

```

```

286         return drawtext
287
288     def _select_encoder(self) -> Tuple[List[str], str]:
289         """Resolve the ``-c:v ?`` args used for EVERY per-clip trim re-encode this run,
plus the
290         chosen encoder name. Chosen ONCE (not per clip) so all clips share one codec and
the final
291         concat can stay ``-c copy``. NVENC (hardware) is the L4 win on the 4K re-encode;
if the
292         requested NVENC encoder can't run in this ffmpeg build (CI / non-GPU box), fall
back to "
293         libx264 so a misconfig never nukes the reel."""
294         enc = self.encoder
295         if enc in ("h264_nvenc", "hevc_nvenc") and not _probe_encoder_available(enc):
296             logger.warning(
297                 "[compile] encoder %r is not runnable in this ffmpeg build ? falling
back to "
298                 "libx264 (CPU). On the L4 worker this should be available.",
299                 enc,
300             )
301         enc = "libx264"
302         if enc == "libx264":
303             # Unchanged historical CPU settings (byte-compatible with the pre-#521
stitch).
304             return ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"], enc
305             # NVENC (H.264/HEVC): hardware encode. -cq is the CRF-equivalent constant-
quality knob; p4 is
306             # a balanced speed/quality preset. yuv420p keeps the output broadly playable.
307             return (
308                 ["-c:v", enc, "-preset", "p4", "-cq", "23", "-pix_fmt", "yuv420p"],
309                 enc,
310             )
311
312     def _trim_one_segment(
313         self,
314         idx: int,
315         raw_start: Union[int, float, str],
316         raw_end: Union[int, float, str],
317         segments: List[Tuple[float, float]],
318         video_duration: float,
319         tmp_dir: str,
320         raw_video_path: str,
321         watermark_filter: Optional[str],
322         video_codec_args: Optional[List[str]] = None,
323     ) -> Tuple[Optional[str], Optional[Tuple[int, float, float, str]]]:
324         """Trim a single segment into its own clip file (extracted verbatim from the
per-segment body of compile_highlights). Returns (clip_path, skip_record): on
325         success the clip path and None; on a skip (invalid/out-of-range/ffmpeg failure)
326         None and the (idx, start, end, reason) record to record in skipped_segments."""
327         # Normalize timestamps to float seconds ? handles float, int,
328         # plain numeric strings, and colon-separated formats like "MM:SS"
329         start = self._parse_time(raw_start)
330         end = self._parse_time(raw_end)
331         segment_label = (
332             f"Segment #{idx + 1}/{len(segments)} [{start:.2f}s - {end:.2f}s]"
333         )
334
335         # Validate the segment times
336         if start < 0:
337             logger.warning(
338                 f"{segment_label}: start_time is negative ({start:.2f}s). Clamping to
0."
339             )
340             start = 0.0
341
342

```



```

343         if start >= end:
344             logger.error(
345                 f"{segment_label}: start_time ({start:.2f}s) >= end_time ({end:.2f}s).
Skipping invalid segment."
346             )
347             return None, (idx, start, end, "start >= end")
348
349         # Check if segment extends beyond video duration
350         if video_duration > 0:
351             if start >= video_duration:
352                 logger.error(
353                     f"{segment_label}: start_time ({start:.2f}s) is beyond the video
duration ({video_duration:.2f}s). "
354                     f"This segment cannot be extracted ? the video ran out. Skipping."
355                 )
356                 return None, (idx, start, end, "start beyond video end")
357
358             if end > video_duration:
359                 original_end = end
360                 end = video_duration
361                 logger.warning(
362                     f"{segment_label}: end_time ({original_end:.2f}s) exceeds video
duration ({video_duration:.2f}s). "
363                     f"Clamping end_time to {video_duration:.2f}s. The segment may be
truncated."
364                 )
365
366             clip_path = os.path.join(tmp_dir, f"clip_{idx}.mp4")
367
368             # Precise sub-second trim using fast re-encoding to preserve stream headers
369             padded_start = max(0.0, start - self.begin_padding)
370             padded_end = end + self.end_padding
371
372             # Clamp padded_end to video duration if known
373             if video_duration > 0 and padded_end > video_duration:
374                 padded_end = video_duration
375                 logger.info(
376                     f"{segment_label}: Padded end ({end + self.end_padding:.2f}s) exceeds
video duration. "
377                     f"Clamping to {video_duration:.2f}s (end padding partially applied)."
378                 )
379
380             trim_duration = padded_end - padded_start
381             logger.info(
382                 f"Trimming {segment_label} -> [{padded_start:.3f}s - {padded_end:.3f}s]
(duration: {trim_duration:.2f}s)"
383             )
384
385             # Default to libx264 (the historical settings) when a caller invokes
_trim_one_segment
386             # directly without resolving an encoder ? keeps this method's old behaviour
intact.
387             codec_args = (
388                 video_codec_args
389                 if video_codec_args is not None
390                 else ["-c:v", "libx264", "-preset", "superfast", "-crf", "22"]
391             )
392             base_cmd = [
393                 ffmpeg_bin(),
394                 "-y",
395                 "-ss",
396                 str(round(padded_start, 3)),
397                 "-to",
398                 str(round(padded_end, 3)),
399                 "-i",

```

```

400         raw_video_path,
401         *codec_args,
402         "-c:a",
403         "aac",
404         "-b:a",
405         "128k",
406     ]
407     vf_args = ["-vf", watermark_filter] if watermark_filter else []
408
409     # Try with the watermark first; if that encode fails (e.g. a bad font/logo
410     # path), retry once WITHOUT it so a branding misconfig can never nuke the
411     # whole reel. Clips stay codec-identical either way, so concat -c copy holds.
412     attempts = [vf_args, []] if vf_args else [[]]
413     produced = False
414     last_exc: Optional[subprocess.CallProcessError] = None
415     cmd = base_cmd + [clip_path]
416     timeout = ffmpeg_timeout_sec()
417     for attempt_vf in attempts:
418         cmd = base_cmd + attempt_vf + [clip_path]
419         try:
420             subprocess.run(
421                 cmd, capture_output=True, check=True, timeout=timeout
422             )
423         except subprocess.TimeoutExpired:
424             reason = f"ffmpeg timeout after {timeout_label(timeout)}"
425             logger.error(f"{segment_label}: FFmpeg trim timed out.")
426             print(f"  Command: {' '.join(cmd)}")
427             print(f"  {reason}")
428             return None, (idx, start, end, reason)
429         except subprocess.CallProcessError as e:
430             last_exc = e
431             if attempt_vf and vf_args:
432                 _err = (
433                     e.stderr.decode(errors="replace").strip()[-300:]
434                     if getattr(e, "stderr", None)
435                     else ""
436                 )
437                 logger.warning(
438                     f"{segment_label}: watermark encode failed; retrying without it."
439                 )
440                 f"ffmpeg: {_err or '(no stderr captured)'}"
441             )
442             continue
443         break
444     if os.path.exists(clip_path) and os.path.getsize(clip_path) > 0:
445         produced = True
446         if not attempt_vf and vf_args:
447             logger.warning(
448                 f"{segment_label}: encoded WITHOUT the watermark (the watermark
449                 filter failed)."
450             )
451         break
452
453     if produced:
454         return clip_path, None
455     elif last_exc is not None:
456         stderr_msg = last_exc.stderr.decode(errors="replace").strip()
457         logger.error(f"{segment_label}: FFmpeg trim failed.")
458         print(f"  Command: {' '.join(cmd)}")
459         print(
460             f"  FFmpeg stderr: {stderr_msg[-500:]}")
461         ) # Last 500 chars to avoid flooding
462         return None, (idx, start, end, f"ffmpeg error: {stderr_msg[-200:]}")

```

```

463         else:
464             logger.error(
465                 f"{segment_label}: FFmpeg exited successfully but output clip is missing
or empty. Skipping."
466             )
467             return None, (idx, start, end, "empty output clip")
468
469     def compile_highlights(
470         self,
471         raw_video_path: str,
472         segments: List[Tuple[float, float]],
473         output_filename: str,
474     ) -> str:
475         """
476         Extracts selected segments from a raw video and stitches them together
477         into a seamless highlights file, applying end_padding to prevent abrupt endings.
478
479         .. warning::
480             This method uses a two-stage codec contract: per-clip re-encodes must use
481             IDENTICAL encoder settings (codec, preset, CRF, pixel format) so the final
482             concat can stream-copy (`-c copy`) without re-encoding. Changing the
483             clip encode settings without updating the concat assumptions will produce
484             broken output.
485         """
486         if not segments:
487             raise ValueError("No highlight segments provided to compile.")
488
489         # Never stitch the same footage twice: collapse overlapping/duplicate ranges
490         # (e.g. accumulated re-run segments or chunk-overlap dupes) before trimming.
491         original_count = len(segments)
492         segments = self._merge_overlapping(segments)
493         if len(segments) < original_count:
494             logger.info(
495                 "Merged %d overlapping/duplicate segment(s) before stitching: %d ->
%d.",
496                 original_count - len(segments),
497                 original_count,
498                 len(segments),
499             )
500         if not segments:
501             raise ValueError("No valid (start < end) highlight segments after merge.")
502
503         if not os.path.exists(raw_video_path):
504             raise FileNotFoundError(
505                 f"[COMPILER ERROR] Source video file not found: {raw_video_path}"
506             )
507
508         # Defense-in-depth: never let the output name escape output_dir (the API also
509         # validates this at the request boundary).
510         output_path = os.path.join(self.output_dir, os.path.basename(output_filename))
511
512         # Probe video duration so we can detect out-of-range segments
513         video_duration = self._get_video_duration(raw_video_path)
514         if video_duration > 0:
515             logger.info(f"Source video duration: {video_duration:.2f}s")
516         else:
517             logger.warning(
518                 "Video duration unknown. Cannot validate segment boundaries against
video length."
519             )
520
521         # We will create temporary clips and stitch them using FFmpeg concat demuxer
522         temp_clips = []
523         skipped_segments = []
524

```

```

525         # Create a temp directory within output_dir
526         with tempfile.TemporaryDirectory(dir=self.output_dir) as tmp_dir:
527             logger.info(
528                 f"Trimming {len(segments)} segments
(begin_padding={self.begin_padding}s, end_padding={self.end_padding}s)..."
529             )
530             # Build the (optional) watermark filtergraph once; it is applied during
every
531             # per-clip re-encode so the mark is baked into each clip identically.
532             watermark_filter = self._watermark_filter(tmp_dir)
533             if watermark_filter:
534                 logger.info("[watermark] applying branding overlay to each clip.")
535             # Resolve the video encoder ONCE (probe NVENC / fall back) so every clip
shares one codec
536             # and the final concat stays -c copy (#521).
537             video_codec_args, chosen_encoder = self._select_encoder()
538             if chosen_encoder != "libx264":
539                 logger.info("[compile] encoding clips with %s (hardware).",
chosen_encoder)
540             for idx, (raw_start, raw_end) in enumerate(segments):
541                 clip_path, skip_record = self._trim_one_segment(
542                     idx,
543                     raw_start,
544                     raw_end,
545                     segments,
546                     video_duration,
547                     tmp_dir,
548                     raw_video_path,
549                     watermark_filter,
550                     video_codec_args,
551                 )
552                 if clip_path is not None:
553                     temp_clips.append(clip_path)
554                 if skip_record is not None:
555                     skipped_segments.append(skip_record)
556
557             # Report summary of skipped segments
558             if skipped_segments:
559                 logger.info(
560                     f"\n[SUMMARY] {len(skipped_segments)} out of {len(segments)}
segments were SKIPPED:"
561                 )
562                 for seg_idx, seg_start, seg_end, reason in skipped_segments:
563                     logger.info(
564                         f"    - Segment #{seg_idx + 1} [{seg_start:.2f}s -
{seg_end:.2f}s]: {reason}"
565                     )
566                 logger.info("")
567
568             if not temp_clips:
569                 raise RuntimeError(
570                     f"[COMPILER ERROR] Failed to generate ANY valid highlight clips. "
571                     f"All {len(segments)} segments were skipped. "
572                     f"Check the errors above for details. "
573                     f"Common causes: video file is corrupted, segments are out of range,
or ffmpeg is not installed."
574                 )
575
576             logger.info(
577                 f"Successfully trimmed {len(temp_clips)}/{len(segments)} clips.
Preparing to stitch..."
578             )
579
580             # Create a list file for FFmpeg concat demuxer
581             list_file_path = os.path.join(tmp_dir, "clips.txt")

```

```

582         # encoding pinned: a non-ASCII clip path would crash on Windows (cp1252
default);
583         # ffmpeg's concat demuxer reads the list as UTF-8 on every OS.
584         with open(list_file_path, "w", encoding="utf-8") as f:
585             for clip in temp_clips:
586                 # The concat demuxer resolves relative entries against the list
file's
587                 # directory. Use absolute paths so a temp-dir clip does not become
588                 # output/tmp.../output/tmp.../clip_0.mp4 on Windows.
589                 safe_path = Path(clip).resolve().as_posix().replace("'", "'\\'")
590                 f.write(f"file '{safe_path}'\n")
591
592         logger.info("Stitching clips together...")
593         # Stitch them together without re-encoding (since they were encoded with
identical settings)
594         concat_cmd = [
595             ffmpeg_bin(),
596             "-y",
597             "-f",
598             "concat",
599             "-safe",
600             "0",
601             "-i",
602             list_file_path,
603             "-c",
604             "copy",
605             output_path,
606         ]
607
608         try:
609             timeout = ffmpeg_timeout_sec()
610             subprocess.run(
611                 concat_cmd, capture_output=True, check=True, timeout=timeout
612             )
613
614             # Final verification
615             if not os.path.exists(output_path) or os.path.getsize(output_path) == 0:
616                 raise RuntimeError(
617                     f"[COMPILER ERROR] FFmpeg concatenation appeared to succeed, "
618                     f"but output file is missing or empty: {output_path}"
619                 )
620
621             output_size_mb = os.path.getsize(output_path) / (1024 * 1024)
622             logger.info(
623                 f"Dynamic highlight reel saved: {output_path} ({output_size_mb:.1f}
MB)"
624             )
625             return output_path
626         except subprocess.CalledProcessError as e:
627             stderr_msg = e.stderr.decode(errors="replace").strip()
628             raise RuntimeError(
629                 f"[COMPILER ERROR] FFmpeg concatenation failed.\n"
630                 f"Command: {' '.join(concat_cmd)}\n"
631                 f"FFmpeg stderr: {stderr_msg[-500:]}"
632             ) from e
633         except subprocess.TimeoutExpired as e:
634             raise RuntimeError(
635                 f"[COMPILER ERROR] FFmpeg concatenation timed out after
{timeout_label(timeout)}.\n"
636                 f"Command: {' '.join(concat_cmd)}"
637             ) from e
638

```

The new `compiler.py` (with #527's probe fix at lines 30-78) is **identical** to my base in every region I edited ? `\_\_init\_\_`, `\_watermark\_filter`, `\_trim\_one\_segment`, `compile\_highlights` are unchanged. So my exact edits re-apply cleanly. Let me re-apply all 6 compiler edits, then commit immediately.

168. user (2026-07-08T07:04:23.192Z)

The file C:\kg-wt-sp\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

169. user (2026-07-08T07:04:30.516Z)

The file C:\kg-wt-sp\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

170. user (2026-07-08T07:04:50.697Z)

The file C:\kg-wt-sp\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

171. user (2026-07-08T07:05:05.295Z)

The file C:\kg-wt-sp\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

172. user (2026-07-08T07:05:33.506Z)

The file C:\kg-wt-sp\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

173. user (2026-07-08T07:05:42.619Z)

The file C:\kg-wt-sp\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

174. assistant (2026-07-08T07:05:47.653Z)

compiler.py fully re-applied. Now `models.py` (Read required for the new worktree path) and `orchestration.py` (touched by #528 ? need to see the current `\_stitch\_reel`).

175. user (2026-07-08T07:05:48.684Z)

```
366         min_rally_duration: float = Field(default=0.0, ge=0.0)
367         # Video encoder for the per-clip trim re-encode (the branding watermark forces a re-
368         # concat stays -c copy). Default ``libx264`` (CPU, works everywhere incl. CI).
369         # ``h264_nvenc`` /
370         # ``hevc_nvenc`` are the L4 HARDWARE encoders ? dramatically faster on the 4K re-
371         # encode that
372         # dominated the 2026-07-08 CUJ stitch (#521). The compiler probes the requested
373         # encoder ONCE and
374         # falls back to libx264 if this ffmpeg build can't run it, so a non-GPU box never
375         # breaks. The
376         # cloud compile-dispatch forwards ``deployment.cloud_stitch_encoder`` onto this for
377         # the L4 run;
378         # the control-plane's own local stitch keeps libx264.
379         encoder: Literal["libx264", "h264_nvenc", "hevc_nvenc"] = "libx264"
380         watermark: WatermarkConfig = Field(default_factory=WatermarkConfig)
```

176. user (2026-07-08T07:05:49.172Z)

```

1206         returns the ``{status, filepath, download_url, video_id, reel_uri}`` the SPA polls
for. Returns a
1207         ``status='failed'`` result on a stitch error (never raises for an expected ffmpeg
failure).
1208
1209         ``require_mirror`` (the L4-worker caller sets True): on the worker the bucket mirror
is the ONLY
1210         delivery ? the worker's local reel is discarded when the Job exits, so a mirror-
upload failure or a
1211         missing bucket destination MUST fail the stitch (else the dispatcher reports success
for a reel the
1212         SPA can never download ? 404 on a "success" job). The control-plane caller keeps the
default False:
1213         its local reel is a valid fallback, so a mirror miss stays best-effort.""
1214         import backend.main as _main # call-time seam: VideoCompiler stays patchable on
backend.main
1215
1216         output_dir = (
1217             getattr(getattr(cfg, "output", None), "output_dir", "output") or "output"
1218         )
1219         # Localize the branding watermark to the UI locale the SPA rendered in (recorded on
the job);
1220         # unknown/absent locale keeps the configured default unchanged.
1221         localized_watermark = localize_watermark(stitching.watermark, locale)
1222         compiler = _main.VideoCompiler(
1223             output_dir,
1224             begin_padding=stitching.begin_padding,
1225             end_padding=stitching.end_padding,
1226             watermark=localized_watermark,
1227             encoder=getattr(stitching, "encoder", "libx264"),

```

177. assistant (2026-07-08T07:05:58.526Z)

Both anchors intact (`#528` added a `require\_mirror` param but not to the `VideoCompiler` construction). Re-applying models.py and orchestration.py edits:

178. user (2026-07-08T07:06:04.233Z)

The file C:\kg-wt-sp\backend\config\models.py has been updated successfully. (file state is current in your context ? no need to Read it back)

179. user (2026-07-08T07:06:08.597Z)

The file C:\kg-wt-sp\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

180. assistant (2026-07-08T07:06:32.589Z)

All code re-applied. Given I already lost this work once, let me **commit** the implementation immediately** to protect it (syntax-check first, verify branch, stage explicit paths only ? never ``-A``), then re-apply tests + docs.

181. user (2026-07-08T07:06:37.287Z)

```

===syntax check===
SYNTAX OK

```

```

===branch safety re-verify===
perf/521-single-pass-stitch
65ee4c6 fix(serving): probe NVENC at 256x256, not 64x64, so it isn't falsely rejected
(#521/#525) (#527)

```

```

===targeted compiler/watermark tests===

```

[100%]

33 passed in 0.72s

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

182. assistant (2026-07-08T07:07:04.613Z)

Syntax clean, branch correct, 33 tests green. Committing the implementation now to protect it (explicit paths only):

183. user (2026-07-08T07:07:11.142Z)

===verify commit is mine and only mine===

650a183 perf(stitch): single-pass concat-filter stitch for a cheaper watermark (#521)

184. assistant (2026-07-08T07:07:22.823Z)

Implementation committed (`650a183`), and it's the only commit above master ? protected. Now re-apply tests + docs. Let me read the (#527/#528-modified) regions I need to anchor into.

185. user (2026-07-08T07:07:25.585Z)

===tail of test_compiler.py (see #527 additions + true EOF)===

dimensions (>=256px per axis). The original 64x64 is BELOW that minimum, so on real GPUs the probe FAILS ("at least one of its streams received no packets") even though NVENC works fine at

real resolutions ? making _select_encoder silently fall back to libx264 and defeating the NVENC-on-L4 re-encode win of #521/#525. Guard the geometry against regressing to a sub-minimum

```
size."""
cmd = _capture_probe_cmd(monkeypatch)
source = cmd[cmd.index("-i") + 1]
match = re.search(r"s=(\d+)\x(\d+)", source)
assert match, f"probe lavfi source has no s=WxH size spec: {source!r}"
width, height = int(match.group(1)), int(match.group(2))
assert width >= 256 and height >= 256, (
    f"probe resolution {width}x{height} is below NVENC's minimum; must be >=256x256 so the "
    "probe reflects real NVENC runnability instead of rejecting an undersized frame
(#521/#525)"
)
```

```
def test_probe_encoder_builds_single_frame_null_encode(monkeypatch):
    """Sanity: the probe stays a real 1-frame test encode to null with the requested codec (not
an
``-encoders`` capability grep), so it still distinguishes 'NVENC compiled-in but unusable'
from
genuinely usable."""
cmd = _capture_probe_cmd(monkeypatch)
assert cmd[cmd.index("-c:v") + 1] == "h264_nvenc" # probes the requested encoder
assert cmd[cmd.index("-frames:v") + 1] == "1" # exactly one frame
assert "null" in cmd # encodes to the null muxer, discarding output
assert "-encoders" not in cmd # a real encode, not a capability grep
```

```
def test_probe_encoder_returns_false_on_ffmpeg_failure(monkeypatch):
    """A non-zero ffmpeg exit (the encoder can't run on this box) ? unavailable ? caller falls
back
to libx264. Previously untested: the existing encoder tests monkeypatch the whole probe
away."""
from backend.pipeline.stitcher import compiler as comp_mod

monkeypatch.setattr(comp_mod, "_ENCODER_AVAILABLE", {})
monkeypatch.setattr(comp_mod.subprocess, "run", lambda *a, **k: MagicMock(returncode=1))
assert comp_mod._probe_encoder_available("h264_nvenc") is False
```



```
=== #527 probe test present? ===
5:import re
250:def test_probe_encoder_uses_safe_resolution(monkeypatch):
Shell cwd was reset to C:\Users\avidu\Projects\khesutra-guru\badminton-highlight-indexer
```

186. assistant (2026-07-08T07:07:39.857Z)

#527 appended three probe tests (+`import re`) at the end of `test\_compiler.py`. My single-pass tests must append after the true EOF. Let me Read the end via the Read tool (required before editing this path), then append.

187. user (2026-07-08T07:07:40.405Z)

```
280         """A non-zero ffmpeg exit (the encoder can't run on this box) ? unavailable ? caller
falls back
281         to libx264. Previously untested: the existing encoder tests monkeypatch the whole
probe away."""
282         from backend.pipeline.stitcher import compiler as comp_mod
283
284         monkeypatch.setattr(comp_mod, "_ENCODER_AVAILABLE", {})
285         monkeypatch.setattr(comp_mod.subprocess, "run", lambda *a, **k:
MagicMock(returncode=1))
286         assert comp_mod._probe_encoder_available("h264_nvenc") is False
287
```

188. user (2026-07-08T07:08:17.082Z)

The file C:\kg-wt-sp\tests\test_compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

189. assistant (2026-07-08T07:08:22.946Z)

Now the dispatch test (#528 modified it). Let me read the nvenc-spec assertion region to re-add my one-line check.

190. user (2026-07-08T07:08:25.206Z)

```
123:def test_compile_dispatch_stages_resolved_spec_with_nvenc(compile_env):
139:     assert fake.spec["stitching"]["encoder"] == "h264_nvenc"
Shell cwd was reset to C:\Users\avidu\Projects\khesutra-guru\badminton-highlight-indexer
```

191. user (2026-07-08T07:08:30.874Z)

```
137         assert fake.spec["time_pairs"] == [[2.0, 6.0], [10.0, 14.0], [20.0, 22.0]]
138         # the L4 stitch uses NVENC (forwarded from deployment.cloud_stitch_encoder)
139         assert fake.spec["stitching"]["encoder"] == "h264_nvenc"
140         # the argv is a `compile` subcommand carrying the spec + done marker
141         argv = fake.calls[0]
142         assert argv[0] == "compile"
```

192. user (2026-07-08T07:08:39.316Z)

The file C:\kg-wt-sp\tests\test_compile_dispatch.py has been updated successfully. (file state is current in your context ? no need to Read it back)

193. assistant (2026-07-08T07:08:48.061Z)

Now the doc. #527/#528 added 8 lines to `PHASE\_PLACEMENT.md` ? let me read the \$3 area to see if my anchor (the follow-up note) survived and how the NVENC section changed.

194. user (2026-07-08T07:08:49.049Z)

100

```

101     **Fallback:** if compile is placed on the L4 but the box isn't actually cloud-wired (no
output bucket
102     / Cloud Run env), `_dispatch_compile_remote` logs and stitches in-process ? a slow reel
beats no reel.
103
104     ---
105
106     ## 3. NVENC on the L4 ? and the D5/D16(b) parity note
107
108     The dominant term in the CUJ stitch was the **per-clip 4K re-encode** the branding
watermark forces
109     (~67 s/clip on 2 vCPU with libx264). The L4 has a hardware H.264/HEVC encoder, so #521
adds an encoder
110     knob:
111
112     - `StitchingConfig.encoder` ? `libx264` (default, CPU, works everywhere incl. CI) |
`h264_nvenc` |
113     `hevc_nvenc`. `VideoCompiler` **probes** the requested encoder once per run (a 1-frame
lavfi test
114     encode) and **falls back to libx264** if this ffmpeg build can't run it, so a
misconfig / non-GPU
115     box never breaks. All clips share the one chosen encoder ? the concat stays `-c copy`.
116     - `deployment.cloud_stitch_encoder` (default `h264_nvenc`) ? what the compile-dispatch
forwards onto
117     the worker's `stitching.encoder` for the L4 run. Set it to `libx264` to run the L4
stitch on CPU.
118     - The control-plane's **own** local stitch always uses `stitching.encoder` (libx264 by
default) ? the
119     CP has no GPU.
120
121     **Parity note (honest):** [README.md](./README.md) **D5/D16(b)** kept stitch **encode =
libx264** to
122     protect a byte-exact stitch parity guarantee, and explicitly **deferred NVENC encode
behind a config
123     knob, "revisit only if M7 cost data justifies relaxing parity." ** The 2026-07-08 CUJ
(~13 min stitch,
124     a hard scaling wall past ~5 concurrent users) **is** that cost data. #521 lands exactly
the deferred
125     shape ? NVENC **behind a config knob**, libx264 the default, config-reversible per
phase. Enabling
126     NVENC relaxes byte-identity **at the reel ENCODE layer only** (a cosmetic, shareable
artifact); the
127     **WINDOW/rally-level parity** (the analytical output) is untouched ? detection and
windowing are
128     unchanged. If bit-exact reels are ever required, set `cloud_stitch_encoder=libx264`.
129
130     > Follow-up (tracked, not in #521's first cut): a **pre-rendered / overlay-filter
watermark** so the
131     > stitch isn't a full re-encode at all (the deeper win beyond swapping the encoder).
132
133     ---
134
135     ## 4. Validation placement + the "nobody validates" guard
136
137     `phase_placement.validation` moves the CPU-bound validator suite off the 2-vCPU control-
plane too,
138     via the existing `--skip-validation` lever (#513):
139
140     - **`control_plane`** (default): the control-plane runs the validators and is the
authoritative gate;
141     the worker **skips** its redundant re-validation (`skip_validation=True`) ? today's
#512/#513
142     behaviour, byte-identical.
143     - **`cloud_run_l4`**: when detection **actually** dispatches to the L4, the control-
plane **delegates**

```

144 ? it does NOT run the validator suite (`\_execute\_processing` skips it), and the dispatch carries

195. assistant (2026-07-08T07:09:06.337Z)

The §3 follow-up blockquote is intact ? my anchor works. Since ****#527 already fixed**** the NVENC probe bug I'd flagged, I'll re-apply my §3.1 with that note updated from "filed" to "fixed in #527". Re-applying:

196. user (2026-07-08T07:09:31.400Z)

The file C:\kg-wt-sp\docs\COMPUTE_DECOUPLED_SERVING\PHASE_PLACEMENT.md has been updated successfully. (file state is current in your context ? no need to Read it back)

197. assistant (2026-07-08T07:09:36.933Z)

Now the §5 Files & tests row. Let me read §5 (it may have gained a #527 row):

198. user (2026-07-08T07:09:39.296Z)

```
223:## 5. Files & tests
232:| NVENC encoder | `backend/pipeline/stitcher/compiler.py` (`encoder`, `_select_encoder`,
`_probe_encoder_available`) | `tests/test_compiler.py` |
Shell cwd was reset to C:\Users\avidu\Projects\khesutra-guru\badminton-highlight-indexer
```

199. user (2026-07-08T07:09:42.400Z)

```
232 | NVENC encoder | `backend/pipeline/stitcher/compiler.py` (`encoder`, `_select_encoder`,
`_probe_encoder_available`) | `tests/test_compiler.py` |
233 | validation delegation | `backend/orchestration.py` (`maybe_dispatch_remote` skip
flag, `_execute_processing`) | `tests/test_api_cloud_dispatch.py` |
```

200. user (2026-07-08T07:09:49.662Z)

The file C:\kg-wt-sp\docs\COMPUTE_DECOUPLED_SERVING\PHASE_PLACEMENT.md has been updated successfully. (file state is current in your context ? no need to Read it back)

201. assistant (2026-07-08T07:09:54.469Z)

All re-applied. Let me run the targeted tests, then fold tests+docs into the commit.

202. user (2026-07-08T07:10:05.723Z)

```
..... [ 82%]
..... [100%]
87 passed in 6.37s
Shell cwd was reset to C:\Users\avidu\Projects\khesutra-guru\badminton-highlight-indexer
```

203. assistant (2026-07-08T07:10:14.805Z)

87 pass (my single-pass tests + #527 probe + #528 dispatch tests all green). Let me fold tests + docs into the commit (explicit paths), then verify the worktree state.

204. user (2026-07-08T07:10:18.506Z)

```
===status before staging===
M docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md
M tests/test_compile_dispatch.py
M tests/test_compiler.py

===amended; files in commit===
8f4e474 perf(stitch): single-pass concat-filter stitch for a cheaper watermark (#521)
backend/config/models.py | 10 +
```

```

backend/orchestration.py | 1 +
backend/pipeline/stitcher/compiler.py | 259 ++++++-----
docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md | 68 +++++-
tests/test_compile_dispatch.py | 2 +
tests/test_compiler.py | 198 ++++++-----
6 files changed, 521 insertions(+), 17 deletions(-)

```

===only my commit above master?===

8f4e474 perf(stitch): single-pass concat-filter stitch for a cheaper watermark (#521)

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

205. assistant (2026-07-08T07:10:49.118Z)

Commit `8f4e474` now holds all 6 files and is the only commit above master ? the work is durably protected. Now the full pre-LGTM gate. Let me run ruff + mypy, and re-verify the re-applied code end-to-end with real ffmpeg (updating the verify script to point at the new worktree, with a robust luma check this time).

206. assistant (2026-07-08T07:11:04.671Z)

"""Re-verify the re-applied single-pass code end-to-end vs REAL ffmpeg (new worktree C:/kg-wt-sp)."""

```

import os
import subprocess
import sys

```

```

WT = r"C:/kg-wt-sp"
sys.path.insert(0, WT)
SP = os.path.dirname(os.path.abspath(__file__))
SRC = os.path.join(SP, "src_4k_hevc10.mp4")
OUT = os.path.join(SP, "verify_out2")
os.makedirs(OUT, exist_ok=True)
from backend.pipeline.stitcher.compiler import VideoCompiler # noqa: E402

```

```

SEG = [(3.0, 7.0), (12.0, 16.0), (20.0, 24.0)]
WM = {"enabled": True, "text": "Generated by Khelsutra.guru"}
exp = sum((e + 1.0) - max(0.0, s - 0.5) for s, e in SEG) # VideoCompiler defaults 0.5/1.0

```

```

def dur(p):
    o = subprocess.run(["ffprobe", "-v", "error", "-show_entries", "format=duration",
                        "-of", "default=nk=1:nw=1", p], capture_output=True, text=True)
    return float(o.stdout.strip())

```

```

def streams(p):
    return subprocess.run(["ffprobe", "-v", "error", "-show_entries",
                           "stream=codec_type,codec_name,width,height", "-of", "csv=p=0", p],
                           capture_output=True, text=True).stdout.replace("\n", " ")

```

```

def corner_mean(reel, t):
    """Mean gray of the bottom-right 700x180 corner at time t (robust raw-bytes method)."""
    p = subprocess.run(["ffmpeg", "-v", "error", "-ss", str(t), "-i", reel, "-frames:v", "1",
                        "-vf", "crop=700:180:in_w-700:in_h-180,format=gray", "-f", "rawvideo",
                        "-"],
                        capture_output=True)
    d = p.stdout
    return (sum(d) / len(d)) if d else -1.0

```

```

pc = VideoCompiler(OUT, watermark=WM, single_pass=False).compile_highlights(SRC, SEG, "pc.mp4")
sp = VideoCompiler(OUT, watermark=WM, single_pass=True).compile_highlights(SRC, SEG, "sp.mp4")
spn = VideoCompiler(OUT, watermark={"enabled": False}, single_pass=True).compile_highlights(SRC,
SEG, "spn.mp4")

```

```

print(f"expected reel ~{exp:.2f}s")
print(f"per-clip      : {dur(pc):.2f}s [{streams(pc)}]")
print(f"single-pass   : {dur(sp):.2f}s [{streams(sp)}]")
print(f"single nowm    : {dur(spn):.2f}s [{streams(spn)}]")

ok = True
def chk(n, c):
    global ok; ok = ok and c; print(f" [{ 'PASS' if c else 'FAIL' }] {n}")

chk("single-pass frame-accurate", abs(dur(sp) - exp) < 0.6)
chk("single-pass == per-clip duration", abs(dur(sp) - dur(pc)) < 0.6)
chk("single-pass 4K video+audio", "3840" in streams(sp) and "audio" in streams(sp))
m_wm, m_no = corner_mean(sp, 3.0), corner_mean(spn, 3.0)
chk(f"watermark rendered (corner mean wm={m_wm:.1f} vs nowm={m_no:.1f})", abs(m_wm - m_no) >
2.0)
print("RESULT:", "ALL PASS" if ok else "SOME FAILED")
sys.exit(0 if ok else 1)

```

207. user (2026-07-08T07:11:41.906Z)

```

===ruff===
All checks passed!
===mypy===
backend\pipeline\validators\base.py:123: note: By default the bodies of untyped functions are
not checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 150 source files
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

208. user (2026-07-08T07:11:41.951Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\scratchpad\verify_sp2.py (file state is current in your context ? no need to Read it back)

209. assistant (2026-07-08T07:11:50.500Z)

ruff + mypy clean (150 files). Let me start the full test suite in the background and run the real-ffmpeg re-verification in the foreground.

210. user (2026-07-08T07:11:52.075Z)

Command running in background with ID: b0kabecar. Output is being written to: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\tasks\b0kabecar.output. You will be notified when it completes. To check interim output, use Read on that file path.

211. user (2026-07-08T07:11:54.357Z)

Command running in background with ID: b3dq9wch9. Output is being written to: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\tasks\b3dq9wch9.output. You will be notified when it completes. To check interim output, use Read on that file path.

212. assistant (2026-07-08T07:12:00.634Z)

Both running. Let me wait for the real-ffmpeg verification (the quicker one).

213. user (2026-07-08T07:13:31.430Z)

```

=== DONE ===
expected reel ~16.50s

```

```
per-clip      : 16.52s [h264,video,3840,2160 aac,audio ]
single-pass   : 16.50s [h264,video,3840,2160 aac,audio ]
single nowm   : 16.50s [h264,video,3840,2160 aac,audio ]
[PASS] single-pass frame-accurate
[PASS] single-pass == per-clip duration
[PASS] single-pass 4K video+audio
[PASS] watermark rendered (corner mean wm=145.2 vs nowm=177.0)
RESULT: ALL PASS
```

214. assistant (2026-07-08T07:13:54.261Z)

****Re-verification ALL PASS**** ? the re-applied code works end-to-end with real ffmpeg: frame-accurate (16.50s), 4K video+audio, watermark rendered (corner mean shifts 177?145 from the box+text).

Since ****#527** already merged my flagged NVENC-probe fix**, that task chip is now stale ? let me dismiss it, then wait for the full suite.

215. user (2026-07-08T07:13:54.630Z)

Task task_8a2b7d23 was already started by the user ? it's no longer pending and can't be withdrawn. Nothing was changed.

216. assistant (2026-07-08T07:14:08.531Z)

The task was already started by the user (that's how #527 got made and merged) ? so it's done, not dismissable. Good. Let me wait for the full test suite.

217. user (2026-07-08T07:17:28.296Z)

```
=== DONE ===
Required test coverage of 70% reached. Total coverage: 85.87%
1889 passed, 8 skipped, 25 warnings in 322.57s (0:05:22)
RESULT: PASS - all tests green
```

218. assistant (2026-07-08T07:17:54.456Z)

```
**Full pre-LGTM gate is GREEN:**
- ? Full suite: **1889 passed, 0 failed**, 8 skipped (the 3 failures in the earlier run were artifacts of the hijacked/corrupt worktree ? gone in the clean one).
- ? Coverage **85.87%** (?70 floor; my tests *raised* it).
- ? ruff clean, mypy clean (150 files).
- ? Real-ffmpeg e2e: frame-accurate, valid 4K+audio, watermark rendered.
```

The task asks for "one PR referencing #521", so let me push and open it. First re-sync and re-verify branch safety (the environment is volatile).

219. user (2026-07-08T07:18:01.359Z)

```
===origin/master tip===
304af7b feat(serving): L4-primary topology design + trust submit-time video_id (#524) (#530)
===my branch: only my commit above master?===
8f4e474 perf(stitch): single-pass concat-filter stitch for a cheaper watermark (#521)
===branch===
perf/521-single-pass-stitch
===push===
remote: Create a pull request for 'perf/521-single-pass-stitch' on GitHub by visiting:
remote:      https://github.com/Khelsutra/badminton-highlight-indexer/pull/new/perf/521-single-pass-stitch
remote:
branch 'perf/521-single-pass-stitch' set up to track 'origin/perf/521-single-pass-stitch'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
```

```
* [new branch]      perf/521-single-pass-stitch -> perf/521-single-pass-stitch
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

220. assistant (2026-07-08T07:18:31.689Z)

Branch pushed and safe on the remote (durable now). origin/master advanced again to `304af7b` (#524/#530) ? my single clean commit sits on a recent ancestor, no problem for the PR. Let me open the PR referencing #521, with the measured results and verification.

221. assistant (2026-07-08T07:18:49.295Z)

Part 3 of #521 (follow-up to the compile-on-L4 + NVENC work merged in #525): make the branding watermark cheaper so the per-clip stitch isn't a full 4K re-encode. Builds on merged #521/#525 and the NVENC-probe fix #527.

The finding (measured, not assumed)

A burned-in watermark covers **every output frame**, so every output frame **must** be encoded ? that encode is an **irreducible floor**. A stream-copy-only stitch (no watermark) is **~1.1 s** vs **~73 s** with the watermark on a real 4K 10-bit HEVC clip: **~98% of the stitch cost IS the watermark re-encode**. So the levers are (a) a faster encoder (NVENC ? already shipped in #521) and (b) removing **per-clip overhead**. This PR is (b).

Measured ? real 4K 10-bit HEVC clip (3840×2160, 8 segments, RTX 2070S, ffmpeg 8.1, warm)

```
| strategy | libx264 | h264_nvenc | frame-accurate? |
|---|---|---|---|
| Baseline (per-clip re-encode + `drawtext` ? concat `-c copy`) | 73.1 s | 34.7 s | ? reel 40.0 s |
| Option A (pre-rendered PNG overlay) | 59.5 s (?19 %) | 36.0 s (**+4 %**) | ? |
| Option B (stream-copy trims ? concat re-encode + wm) | 65.3 s (?11 %) | 28.9 s (?17 %) | ?
**reel 48.1 s** |
| **Option C ? single-pass (chosen)** | **67.4 s (?14 %)** | **34.1 s (?11 %)** | ? reel 40.0 s |
|
| _floor: stream-copy, no watermark_ | 1.1 s | 1.2 s | ? |
```

Chose Option C: one ffmpeg pass ? precise `-ss/-t` inputs ? concat **filter** ? watermark **once** ? a single encode.

- Unlike the brief's **Option B** (stream-copy trims ? concat re-encode), the concat **filter** keeps cuts **frame-accurate**; Option B's stream-copy snapped to keyframes (48.1 s vs the correct 40.0 s reel ? a **content** regression, not a cosmetic encode relaxation).
- Unlike **Option A**, it helps **NVENC** too (the L4 primary path; A regresses NVENC ~4 %).
- Also produces **smaller** reels (one continuous encode vs N keyframe-forced GOPs) and writes **no intermediate clips**.

Safety

- **Opt-in** ``stitching.single_pass_stitch`` (default **OFF**). The per-clip path stays the parity-safe default **and** the automatic **fallback**: ``_compile_single_pass`` returns `None` on any failure (ffmpeg error, timeout, segment-count cap, no valid segments) and ``compile_highlights`` transparently re-runs the per-clip path ? a slower reel beats no reel.
- Segments validated up front via the shared ``_resolve_padded_window`` (extracted from ``_trim_one_segment``, behaviour-preserving), so an invalid segment is dropped before the graph.
- Rides the compile-spec to the L4 worker exactly like ``encoder``.
- **Parity (D5/D16(b)):** identical footage (frame-accurate) + identical watermark; only the **encode structure** differs (one continuous encode vs N GOPs) ? within the reel-**ENCODE**-layer relaxation D16(b) already allows under NVENC. ``single_pass_stitch=false`` keeps the old path bit-for-bit.

Verification

- **Full suite:** 1889 passed, 0 failed, 8 skipped; **coverage 85.87 %** (? 70 floor; new tests

```
raise it).
- **ruff + mypy clean** (150 source files).
- **Real ffmpeg end-to-end:** single-pass reel frame-accurate (16.50 s vs per-clip 16.52 s),
valid 4K H.264 + AAC, watermark rendered (verified on the RTX 2070S).
- **Owner-gated:** a confirming run on the **L4 + golden corpus** (per the repo's machine-side
split); the transferable result is the relative deltas + frame-accuracy.
```

Docs: `docs/COMPUTE\_DECOUPLED\_SERVING/PHASE\_PLACEMENT.md` **\$3.1** ? full measured before/after and why A/B were rejected.

? Generated with [Claude Code](https://claude.com/claude-code)

222. user (2026-07-08T07:18:49.344Z)

File created successfully at: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\scratchpad\pr_body.md (file state is current in your context ? no need to Read it back)

223. user (2026-07-08T07:18:57.922Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/531>
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

224. assistant (2026-07-08T07:19:11.991Z)

PR created: **<https://github.com/Khelsutra/badminton-highlight-indexer/pull/531>** . Let me verify it's clean (only my commit, mergeable against the advanced master) and check CI kicked off.

225. user (2026-07-08T07:19:17.006Z)

PR # 531 base: master
mergeable: MERGEABLE UNSTABLE
commits: ['perf(stitch): single-pass concat-filter stitch for a cheaper watermar?']
files: 6 -> ['backend/config/models.py', 'backend/orchestration.py',
'backend/pipeline/stitcher/compiler.py', 'docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md',
'tests/test_compile_dispatch.py', 'tests/test_compiler.py']
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

226. assistant (2026-07-08T07:20:06.358Z)

PR **\$531** is clean: **MERGEABLE**, exactly **1 commit** (mine, no sibling commits absorbed), exactly **6 files** (my intended set). `UNSTABLE` just means CI is still running ? and I've already reproduced every CI gate locally (ruff, mypy, full suite, coverage). No conflicts with the twice-advanced master.

Let me check the project memory dir to record a durable pointer for future sessions (without touching the sibling-owned repo-root `SESSION\_HANDOFF.md`).

227. user (2026-07-08T07:20:08.048Z)

```
total 24
drwxr-xr-x 1 avidu 197609    0 Jul  8 12:18 .
drwxr-xr-x 1 avidu 197609    0 Jul  8 12:11 ..
-rw-r--r-- 2 avidu 197609  995 Jul  8 12:18 MEMORY.md
-rw-r--r-- 1 avidu 197609 2920 Jul  8 11:34 pr-525-compile-on-l4.md
-rw-r--r-- 2 avidu 197609 2664 Jul  8 12:18 pr-527-nvenc-probe-resolution.md
-rw-r--r-- 2 avidu 197609 1709 Jul  8 12:18 pr-528-worker-mirror-required.md
```

228. assistant (2026-07-08T07:20:20.838Z)

Sibling sessions record each PR as a memory file here (pr-525, pr-527, pr-528) ? an active convention. Let me follow it for my PR #531. Reading the index first to append safely:

229. user (2026-07-08T07:20:21.161Z)

```
1      # Memory index
2
3      - [PR #525 ? compile-on-L4 (#521)](pr-525-compile-on-l4.md) ? config-driven
phase?machine placement + L4 NVENC stitch; Part 3 (cheaper watermark) deferred.
4      - [PR #527 ? NVENC probe 64x64?256x256](pr-527-nvenc-probe-resolution.md) ? sub-minimum
probe silently defeated #521's NVENC-on-L4 win; MERGED to master (65ee4c6); STILL awaits owner
live-GPU confirm; monitoring follow-up #529.
5      - [PR #528 ? worker mirror-required](pr-528-worker-mirror-required.md) ? #521 follow-up:
L4-worker stitch now fails on a lost reel mirror-upload (was silent 404-on-success); found by
adversarial review; **MERGED to master** (008a5d4).
6      - **Release baseline (2026-07-08):** master `65ee4c6` carries #521 (compile-on-L4) + its
two follow-up fixes #528 (mirror-required) + #527 (NVENC 256px probe). Owner cut/cutting a
release folding all three. Real L4/NVENC exec still unverified (needs a worker image rebuilt
from this master; deployed `:latest` predated the compile subcommand).
7
```

230. assistant (2026-07-08T07:20:32.771Z)

Let me record PR #531 as a project memory following the established pattern, and add its index pointer:

231. assistant (2026-07-08T07:20:48.399Z)

name: pr-531-single-pass-stitch

description: PR #531 ? #521 Part-3 (cheaper watermark) via an opt-in single-pass concat-filter stitch; measured, default-OFF, awaits owner L4/corpus confirm.

metadata:

type: project

PR [531](https://github.com/Khelsutra/badminton-highlight-indexer/pull/531) ?
`perf(stitch): single-pass concat-filter stitch for a cheaper watermark (#521)`. Branch
`perf/521-single-pass-stitch` off master; the deferred Part-3 of [[pr-525-compile-on-l4]].

What it does: opt-in `stitching.single\_pass\_stitch` (default **OFF**). When on,
`VideoCompiler.\_compile\_single\_pass` stitches the reel in ONE ffmpeg pass ? precise `-ss/-t`
inputs ? `concat` **filter** ? watermark drawn once ? a single encode ? instead of per-clip re-
encode + concat-copy. Falls back to the per-clip path on any failure; shared
`\_resolve\_padded\_window` validates segments up front; `\_SINGLE\_PASS\_MAX\_SEGMENTS` cap guards the
N-input graph. Rides the compile-spec to the L4 like `encoder`.

Key finding (why the brief's framing was corrected): a burned-in watermark re-encodes EVERY
frame ? that encode is an irreducible floor (a no-watermark stream-copy stitch is ~1.1 s vs ~73
s with it; ~98?99 % of stitch cost IS the watermark re-encode). So neither pre-render (Option A)
nor stream-copy-then-concat-re-encode (Option B) removes the floor; the win is per-clip
OVERHEAD. Measured on a real 4K 10-bit HEVC clip (RTX 2070S): single-pass (Option C) **?14 %**
libx264 / ?11 % NVENC, frame-accurate, smaller output. Option B was rejected (stream-copy
snaps to keyframes ? reel 48.1 s vs 40.0 s, a content regression); Option A helps only libx264
(+4 % on NVENC).

Verified: full suite 1889 passed / coverage 85.87 %; ruff+mypy clean; real-ffmpeg e2e
(frame-accurate, 4K+audio, watermark rendered). **Owner-gated:** confirming run on the L4 +
golden corpus (machine-side split). Doc: `docs/COMPUTE\_DECOUPLED\_SERVING/PHASE\_PLACEMENT.md`
§3.1.

Process note (2026-07-08): built in an isolated worktree; a concurrent codex session
HIJACKED the first worktree (checked out `codex/issue-523-sd-reel-review` + reset it), wiping my
UNCOMMITTED changes. Recovered from conversation history, re-applied onto the then-current
master (which had gained #527/#528), and **committed immediately**. Lesson: in this high-
concurrency checkout, commit early ? uncommitted worktree edits are not safe. Related: my
flagged NVENC-probe bug became [[pr-527-nvenc-probe-resolution]] (merged before this PR).

232. user (2026-07-08T07:20:48.468Z)

File created successfully at: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\pr-531-single-pass-stitch.md (file state is current in your context ? no need to Read it back)

233. user (2026-07-08T07:21:00.039Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

234. assistant (2026-07-08T07:21:09.129Z)

Memory recorded. Let me check the PR's CI status to report accurately, then summarize.

235. user (2026-07-08T07:21:12.567Z)

```
Golden fixtures cross-OS (macos-latest) pass    22s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214347
Golden fixtures cross-OS (windows-latest)      pass    53s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214398
Leak scan (PII / attribution guard)            pass    9s      https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214296
Markdown link check (live docs) pass          6s      https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214322
Mypy (lenient baseline) pass                   41s      https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214332
Ruff lint                                       pass    11s      https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214355
Syntax sweep + import smoke (no GPU stack)      pass    25s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214309
Full test suite (offline, mocked)              pending 0      https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214368
Built wheel smoke (install + launch)           pass    1m3s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214344
Dependency license scan (no copyleft in runtime deps) pass    33s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214330
Golden fixtures cross-OS (ubuntu-latest)        pass    26s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28924919234/job/85810214352
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

236. assistant (2026-07-08T07:21:52.855Z)

All CI checks **pass** except the full mocked suite which is still **pending** (running) ? and I've reproduced that locally (1889 passed). Notably the "Syntax sweep + import smoke" check **passed**, confirming the 3 failures in my earlier run were artifacts of the hijacked worktree, not real. The work is done.

Summary

Delivered: [PR #531](https://github.com/Khelsutra/badminton-highlight-indexer/pull/531) ? the deferred **Part 3 of #521**: make the branding watermark cheaper so the per-clip stitch isn't a full 4K re-encode.

What the investigation found (the task's core)

I measured, rather than assumed. A burned-in watermark covers **every frame**, so every frame **must** be encoded ? that encode is an **irreducible floor**. On a real 4K 10-bit HEVC clip, a no-watermark stream-copy stitch is **~1.1s** vs **~73s** with the watermark, i.e. **~98?99%** of stitch cost is the watermark re-encode. So the only levers are a faster encoder (NVENC ?

already shipped in #521) and removing **per-clip overhead**.

I benchmarked the brief's Option A and B plus a third variant, on **both** encoders:

```
| strategy | libx264 | h264_nvenc | frame-accurate |
|---|---|---|---|
| Baseline (current) | 73.1s | 34.7s | ? |
| Option A (pre-rendered overlay) | 59.5s (?19%) | 36.0s (**+4%**) | ? |
| Option B (stream-copy ? concat re-encode) | 65.3s | 28.9s | ? **48.1s reel** |
| **Option C ? single-pass (chosen)** | **67.4s (?14%**) | **34.1s (?11%**) | ? |
```

Chose Option C ? one ffmpeg pass (precise `-ss/-t` inputs ? `concat` `filter` ? watermark once ? single encode). It beats the brief's Option B, which stream-copies to keyframes and produces a **wrong-length reel** (48.1s vs 40.0s ? a content regression, not the cosmetic encode relaxation D16(b) allows), and beats Option A, which regresses NVENC (the L4 primary path). It's **opt-in** (`stitching.single_pass_stitch`, default OFF), with the robust per-clip path as both the parity-safe default and the automatic fallback.

Verified (pre-LGTM gate, all reproduced locally)

- Full suite **1889** passed, 0 failed; coverage **85.87%** (?70 floor, my tests raise it) . ruff + mypy clean.
- **Real-ffmpeg end-to-end**: single-pass reel frame-accurate (16.50s vs per-clip 16.52s), valid 4K H.264+AAC, watermark rendered (visually confirmed).
- CI: all checks green except the full suite (still running).
- Documented in `docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md` **\$3.1** with the measured before/after and A/B rejection rationale.
- **Owner-gated** (per the repo's machine-side split): a confirming run on the L4 + golden corpus ? the transferable result is the relative deltas + frame-accuracy on a synthetic clip.

Two things you should know

1. **A latent #521 bug** I flagged is already fixed and merged (#527):
`_probe_encoder_available` probed NVENC at `64x64`, below NVENC's minimum encode size, so it **falsely** fell back to libx264 on real GPUs ? silently defeating #521's NVENC-on-L4 win. You (via the task chip) got it fixed to `256x256`. My doc note reflects that it's now fixed.
2. **A concurrent session hijacked** my first isolated worktree ? it checked out `codex/issue-523-sd-reel-review` into it and reset, wiping my **uncommitted** changes mid-way. I recovered from context, re-applied onto the (meanwhile-advanced) master, and **committed** immediately this time. The lesson is baked into a project memory: in this many-worktree checkout, commit early. I deliberately did **not** touch the sibling-owned root `SESSION_HANDOFF.md` (active #523 WIP) ? recorded this thread in the project memory dir instead.

The branch is pushed and durable; my local worktree `C:/kg-wt-sp` holds it and can be pruned after merge.

237. user (2026-07-08T07:21:55.763Z)

please make this as default on so that the next release can pick this up

238. user (2026-07-08T07:22:16.138Z)

[Request interrupted by user]

239. user (2026-07-08T07:53:09.545Z)

merged

240. assistant (2026-07-08T07:55:27.844Z)

#531 is merged ? good. Now you want `single_pass_stitch` **default ON** so the next release picks it up. That's a small, focused default-flip on a fresh branch off the latest master. Let me first check my worktree is still intact (hijack risk), re-sync, and confirm nothing else pins the default OFF.

241. user (2026-07-08T07:55:35.897Z)

```
===worktree state===
branch: perf/521-single-pass-stitch

origin/master: 304af7b feat(serving): L4-primary topology design + trust submit-time video_id
(#524) (#530)

===is #531 (single-pass) merged into master?===
NO ? not on master yet

===does config.json pin it?===
config.json: no single_pass key (Pydantic default governs)

===does cloud-serving preset touch it?===
presets.py: no single_pass key

===any test asserting the CONFIG default is False?===
tests/test_compile_dispatch.py:141: assert "single_pass_stitch" in fake.spec["stitching"]
Binary file tests/__pycache__/test_compile_dispatch.cpython-313-pytest-9.0.3.pyc matches
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

242. assistant (2026-07-08T07:55:56.635Z)

My worktree is intact this time. But like before, my local `origin/master` (304af7b) doesn't yet show #531's `single\_pass\_stitch` ? likely the same fetch-lag. You said it's merged, so let me re-sync harder and confirm where #531 landed before branching.

243. user (2026-07-08T07:56:01.935Z)

```
===origin/master now===
304af7b feat(serving): L4-primary topology design + trust submit-time video_id (#524) (#530)
65ee4c6 fix(serving): probe NVENC at 256x256, not 64x64, so it isn't falsely rejected
(#521/#525) (#527)
008a5d4 fix(compile): fail the L4-worker stitch when the reel mirror-upload fails (#521 follow-
up) (#528)
dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521) (#525)
bf97700 docs(runlogs): record live SPA CUJ closing the 2 owner-gated validation-latency cells
($9) (#522)

===single_pass_stitch on origin/master?===
ABSENT on master

===search master history for #531===
351a27e perf(eval): single-pass person detection for fusion_golden (~10x faster extraction)
(#125)
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

244. assistant (2026-07-08T07:56:21.729Z)

A hard `fetch --prune` still shows #531 **absent** from `origin/master` ? the only "single-pass" hit is an unrelated old PR. Your "merged" and my git view disagree, so let me check the authoritative GitHub state for #531 before I branch (this determines whether I flip on master, or just add the default-on commit to the still-open #531).

245. user (2026-07-08T07:56:25.479Z)

```
state: OPEN
mergedAt: None
mergeCommit: None
base: master head: perf/521-single-pass-stitch
mergeable: MERGEABLE CLEAN
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

246. assistant (2026-07-08T07:57:44.451Z)

GitHub's authoritative API says ****#531 is OPEN, not merged**** (``mergedAt: None`, `mergeable: CLEAN` ? CI went green`). So your "merged" and the actual state differ. Since #531 is still open, the cleanest way to make single-pass the default is to add the default-on flip ****to #531 itself**** (keeps it one PR referencing #521, as intended) rather than a second PR. Let me verify my branch is synced with the PR head, then make the change.

247. user (2026-07-08T07:57:47.642Z)

```
===local vs origin PR head===
local: 8f4e474124ac23e1dbe52838d55d669cca5a3b5b
origin: 8f4e474124ac23e1dbe52838d55d669cca5a3b5b
IN SYNC (safe to add a commit)

===current default line in models.py===
384:     single_pass_stitch: bool = False
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

248. assistant (2026-07-08T07:58:29.995Z)

In sync at ``8f4e474``. Now let me make the default-on change ? flip the config default and clarify the comments (keeping the ``VideoCompiler`` constructor default ``False`` as the conservative primitive + fallback, so the ~10 per-clip unit tests that construct it directly stay valid).

249. user (2026-07-08T07:58:34.726Z)

The file `C:\kg-wt-sp\backend\config\models.py` has been updated successfully. (file state is current in your context ? no need to Read it back)

250. user (2026-07-08T07:58:43.118Z)

The file `C:\kg-wt-sp\backend\pipeline\stitcher\compiler.py` has been updated successfully. (file state is current in your context ? no need to Read it back)

251. assistant (2026-07-08T07:58:49.280Z)

Now a config-default test. Let me read the config test section of ``test_watermark.py`` (this worktree path) to place it:

252. user (2026-07-08T07:58:49.856Z)

```
20     # --- config -----
21
22
23     def test_watermark_config_defaults_on_with_brand_text():
24         wm = AppConfig().stitching.watermark
25         assert wm.enabled is True # default-on, with an off switch
26         assert wm.text == "Generated by Khelsutra.guru"
27         assert wm.position == "bottom-right"
28
29
30     def test_watermark_can_be_disabled_via_config():
31         cfg = AppConfig.model_validate({"stitching": {"watermark": {"enabled": False}}})
32         assert cfg.stitching.watermark.enabled is False
33
34
35     def test_font_size_ratio_and_opacity_bounds_enforced():
36         with pytest.raises(ValidationError):
37             WatermarkConfig(opacity=1.5)
```

253. user (2026-07-08T07:58:57.812Z)

The file `C:\kg-wt-sp\tests\test_watermark.py` has been updated successfully. (file state is current in your context ? no need to Read it back)

254. assistant (2026-07-08T07:59:10.199Z)

Now update the doc §3.1 to reflect default-on. Let me read the exact current lines:

255. user (2026-07-08T07:59:12.658Z)

```
132:> **Status:** ? SHIPPED behind `stitching.single_pass_stitch` (default OFF) · one PR
referencing #521.
176:- `StitchingConfig.single_pass_stitch` ? `false` (default). When **true**, `VideoCompiler`
stitches the
178:- **Default OFF = parity-safe.** The per-clip path stays the default *and* the automatic
**fallback**:
```

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

256. user (2026-07-08T07:59:17.723Z)

```
130     ### 3.1 Part-3 ? cheaper watermark: the single-pass stitch (? shipped, measured)
131
132     > **Status:** ? SHIPPED behind `stitching.single_pass_stitch` (default OFF) · one PR
referencing #521.
133     > Was tracked here as "a pre-rendered / overlay-filter watermark so the stitch isn't a
full re-encode
134     > at all." **Investigated + measured on a real 4K clip; the honest finding reshaped the
fix** (below).
135
```

257. user (2026-07-08T07:59:18.451Z)

```
174     **Shape (mirrors the §3 encoder knob ? opt-in, default-safe, self-healing).**
175
176     - `StitchingConfig.single_pass_stitch` ? `false` (default). When **true**,
`VideoCompiler` stitches the
177     reel in one pass (`_compile_single_pass`). Rides the compile-spec onto the L4 like
`encoder` does.
178     - **Default OFF = parity-safe.** The per-clip path stays the default *and* the automatic
**fallback**:
```

```
179     `_compile_single_pass` returns `None` on any failure (ffmpeg error, timeout,
>`_SINGLE_PASS_MAX_SEGMENTS`
180     cap, no valid segments) and `compile_highlights` transparently re-runs the per-clip
path ? a slower
181     reel beats no reel. Segments are validated up front via the shared
`_resolve_padded_window`, so an
182     invalid segment is dropped *before* the graph, never failing the whole encode.
183     - **Parity:** identical footage (frame-accurate) and an identical watermark ? only the
*encode structure*
184     differs (one continuous encode vs N GOPs), which is exactly the reel-ENCODE-layer
relaxation D16(b)
185     already allows under NVENC. WINDOW/rally-level parity is untouched. To force bit-for-
bit the old path,
186     leave `single_pass_stitch=false`.
187
```

258. assistant (2026-07-08T07:59:41.187Z)

Applying the doc §3.1 default-on updates:

259. user (2026-07-08T07:59:44.415Z)

The file C:\kg-wt-sp\docs\COMPUTE_DECOUPLED_SERVING\PHASE_PLACEMENT.md has been updated successfully. (file state is current in your context ? no need to Read it back)

260. user (2026-07-08T07:59:55.713Z)

The file C:\kg-wt-sp\docs\COMPUTE_DECOUPLED_SERVING\PHASE_PLACEMENT.md has been updated

successfully. (file state is current in your context ? no need to Read it back)

261. assistant (2026-07-08T08:00:03.443Z)

Let me update the "Honest limits" note to reflect that default-on raises the priority of the L4 confirmation (while the fallback protects). Let me locate it:

262. user (2026-07-08T08:00:05.299Z)

```
192:**Honest limits (owner-gated, per the repo's machine-side split).** Numbers are from a
**synthetic**
198:`VideoCompiler` NVENC path is correctly enabled. A confirming run on the L4 + golden corpus
stays owner-gated.
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

263. user (2026-07-08T08:00:10.963Z)

The file C:\kg-wt-sp\docs\COMPUTE_DECOUPLED_SERVING\PHASE_PLACEMENT.md has been updated successfully. (file state is current in your context ? no need to Read it back)

264. assistant (2026-07-08T08:00:20.613Z)

All default-on edits done. Let me syntax-check and run the targeted tests (watermark/compiler/dispatch + any config default tests) to confirm the flip breaks nothing:

265. user (2026-07-08T08:01:14.437Z)

SYNTAX OK

```
INFO      backend.pipeline.stitcher.compiler:compiler.py:683 Dynamic highlight reel saved
(single-pass): C:\Users\avidu\AppData\Local\Temp\pytest-of-
avidu\pytest-1643\test_compile_uses_configured_s0\out\reel.mp4 (0.0 MB)
INFO      httpx:_client.py:1025 HTTP Request: POST http://testserver/api/compile "HTTP/1.1 202
Accepted"
INFO      httpx:_client.py:1025 HTTP Request: GET
http://testserver/api/jobs/4b8e39996be74f9e9e1b28a056c2432b "HTTP/1.1 200 OK"
===== warnings summary =====
tests/test_api_endpoints.py: 10 warnings
  C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\site-
packages\google\auth\_default.py:113: UserWarning: Your application has authenticated using end
user credentials from Google Cloud SDK without a quota project. You might receive a "quota
exceeded" or "API not enabled" error. See the following page for troubleshooting:
https://cloud.google.com/docs/authentication/adc-troubleshooting/user-creds.
    warnings.warn(_CLOUD_SDK_CREDENTIALS_WARNING)

-- Docs: https://docs.pytest.org/en/stable/how-to/capture-warnings.html
===== short test summary info =====
FAILED tests/test_api_endpoints.py::test_compile_uses_configured_stitching_paddings
1 failed, 175 passed, 10 warnings in 47.29s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

266. assistant (2026-07-08T08:01:36.165Z)

A test failed: `test\_compile\_uses\_configured\_stitching\_paddings` ? and the log shows it now runs the `**single-pass**` path (config default flipped), so its assertions about per-clip trim commands no longer match. This is a legitimate test that assumed per-clip; I need to fix it properly, not just paper over it. Let me read it.

267. user (2026-07-08T08:01:37.676Z)

```
606:def test_compile_uses_configured_stitching_paddings(env):
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

```

606     def test_compile_uses_configured_stitching_paddings(env):
607         """config.json stitching.{begin,end}_padding must reach the ffmpeg trim boundaries.
608         Regression for the 2026-06-10 audit finding (d): /api/compile constructed
609         VideoCompiler without the paddings, so reels always used the constructor
610         defaults (0.5/1.0) and the config values were silently ignored."""
611         client, db, tmp_path, mp = env
612         cfg = m.global_config
613         cfg.output.output_dir = str(tmp_path / "out")
614         cfg.stitching.begin_padding = 2.0
615         cfg.stitching.end_padding = 3.0
616         cfg.stitching.watermark.enabled = False # keep the trim command minimal
617         db.add_video("vid", str(_video(tmp_path)), "processed", [])
618         sid = db.add_play_segment("vid", 10.0, 14.0)
619
620         calls = []
621
622         def fake_run(cmd, *a, **k):
623             calls.append(cmd)
624             if cmd[0] == "ffprobe":
625                 return MagicMock(stdout="100.0\n", returncode=0)
626             out = cmd[-1]
627             if isinstance(out, str) and out.endswith(".mp4"):
628                 open(out, "wb").write(b"clip")
629             return MagicMock(returncode=0, stdout=b"", stderr=b"")
630
631         with patch.object(subprocess, "run", side_effect=fake_run):
632             # The inline executor runs the stitch (under this patch) during the POST, so the
633             # trim commands are captured here; then the job is already 'done'.
634             r = client.post(
635                 "/api/compile",
636                 json={
637                     "video_id": "vid",
638                     "segment_ids": [sid],
639                     "output_filename": "reel.mp4",
640                 },
641             )
642             job = client.get(f"/api/jobs/{r.json()['job_id']}").json()
643             assert r.status_code == 202
644             assert job["status"] == "done" and job["result"]["status"] == "success"
645
646             trim = next(c for c in calls if c[0] == "ffmpeg" and "-ss" in c)
647             # 10.0 - begin_padding 2.0 and 14.0 + end_padding 3.0; the unwired constructor
648             # defaults (0.5/1.0) would have produced 9.5/15.0 here.
649             assert trim[trim.index("-ss") + 1] == "8.0"
650             assert trim[trim.index("-to") + 1] == "17.0"
651
652
653     def test_parity_job_worker_vs_direct_cli_path(env, tmp_path, monkeypatch):
654         """Item 5 positive: same ProcessRequest through job worker path and direct CLI-style
655         path
656         produce byte-identical DB rows, reports, final play_segments (modulo timestamps).
657         Exercises Failure and empty-run branches.
658         """
659         client, db, tpath, mp = env
660         vid_file = _video(tpath)
661         vid = m.compute_video_id(str(vid_file))
662         mp.setattr(m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
663         mp.setattr(m.segmenter_registry, "get", lambda name: _FakeSeg())
664         # job worker path
665         rj = client.post(
666             "/api/process", json={"video_path": str(vid_file), "segmenter": "x"}

```



```

667         jobj = client.get(f"/api/jobs/{rj.json()['job_id']}").json()
668         # direct CLI-style path (using same processing setup as main CLI; inline executor
makes outcomes identical in test harness)
669         rd = client.post(
670             "/api/process", json={"video_path": str(vid_file), "segmenter": "x"}
671         )
672         jobd = client.get(f"/api/jobs/{rd.json()['job_id']}").json()
673         # byte-identical key outcomes
674         assert jobj["result"]["status"] == jobd["result"]["status"]
675         assert len(db.query_play_segments(vid, {})) == len(db.query_play_segments(vid, {}))

```

269. assistant (2026-07-08T08:02:32.354Z)

Clear: the padding regression test now hits the single-pass default, which expresses the window `[8.0, 17.0]` as input seek `-ss 8.0 -t 9.0` rather than `-ss 8.0 -to 17.0`. The paddings *are* correctly applied ? the test's assertion form was per-clip-specific. Since single-pass is what ships now, I'll update it to assert the default path's form:

270. user (2026-07-08T08:02:34.598Z)

The file C:\kg-wt-sp\tests\test_api_endpoints.py has been updated successfully. (file state is current in your context ? no need to Read it back)

271. assistant (2026-07-08T08:02:44.080Z)

Let me confirm that fix and re-run the affected tests, then launch the full suite (authoritative for a default-flip).

272. user (2026-07-08T08:02:52.936Z)

```

..... [100%]
42 passed in 3.47s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

273. assistant (2026-07-08T08:02:58.381Z)

Let me run the full suite (background, authoritative for a default-flip) and ruff/mypy (foreground).

274. user (2026-07-08T08:02:59.917Z)

Command running in background with ID: bqyazpqzg. Output is being written to: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\tasks\bqyazpqzg.output. You will be notified when it completes. To check interim output, use Read on that file path.

275. user (2026-07-08T08:03:07.849Z)

```

===ruff===
All checks passed!
===mypy===
backend\providers\__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 150 source files
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

276. assistant (2026-07-08T08:03:21.093Z)

ruff + mypy clean. Let me wait for the full suite (authoritative for the default flip ? it'll catch any other test that assumed per-clip).

277. user (2026-07-08T08:06:49.888Z)

=== DONE ===

Required test coverage of 70% reached. Total coverage: 85.87%

1890 passed, 8 skipped, 25 warnings in 222.22s (0:03:42)

RESULT: PASS - all tests green

278. assistant (2026-07-08T08:07:08.805Z)

Full suite green with the default flipped: ****1890 passed, 0 failed**** (the one padding test was the only casualty, now fixed to assert the single-pass form), coverage 85.87%, ruff + mypy clean. Committing the default-flip as a distinct, revertable commit on the #531 branch and pushing:

279. user (2026-07-08T08:07:20.849Z)

branch: perf/521-single-pass-stitch | in sync with PR head: yes

===my commits above master===

b98209f perf(stitch): default single_pass_stitch ON so the next release picks it up (#521)

8f4e474 perf(stitch): single-pass concat-filter stitch for a cheaper watermark (#521)

===push===

To <https://github.com/Khelsutra/badminton-highlight-indexer.git>

8f4e474..b98209f perf/521-single-pass-stitch -> perf/521-single-pass-stitch

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

280. assistant (2026-07-08T08:07:37.724Z)

PR #531 now has 2 commits (feature + default-on), both mine above master. Let me update the PR description to reflect default-on (it currently describes the opt-in version), then update the memory note.

281. user (2026-07-08T08:07:43.931Z)

The file C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\scratchpad\pr_body.md has been updated successfully. (file state is current in your context ? no need to Read it back)

282. user (2026-07-08T08:07:52.120Z)

The file C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\scratchpad\pr_body.md has been updated successfully. (file state is current in your context ? no need to Read it back)

283. assistant (2026-07-08T08:07:57.362Z)

Let me update the verification line to the current suite count and note the flip was verified, then apply the body to the PR:

284. user (2026-07-08T08:08:00.667Z)

The file C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\46eabfd2-0e71-4ff5-a68d-c76f474d4ff3\scratchpad\pr_body.md has been updated successfully. (file state is current in your context ? no need to Read it back)

285. user (2026-07-08T08:08:17.894Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/531>

===PR state + commits===

state: OPEN | mergeable: MERGEABLE UNSTABLE

commits: ['perf(stitch): single-pass concat-filter stitch for a cheaper watermar?',
'perf(stitch): default single_pass_stitch ON so the next release picks?']

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

286. assistant (2026-07-08T08:08:31.862Z)

PR #531 updated ? 2 commits (feature + default-on), OPEN, mergeable (CI re-running). Let me update my memory note to reflect default-on:

287. user (2026-07-08T08:08:34.804Z)

<tool_use_error>String to replace not found in file.
String: description: PR #531 ? #521 Part-3 (cheaper watermark) via an opt-in single-pass concat-filter stitch; measured, default-OFF, awaits owner L4/corpus confirm.
(note: Edit also tried swapping \uXXXX escapes and their characters; neither form matched, so the mismatch is likely elsewhere in old_string. Re-read the file and copy the exact surrounding text.)</tool_use_error>

288. assistant (2026-07-08T08:08:39.249Z)

The memory file content differs from what I expect ? let me re-read it (it may have been touched):

289. user (2026-07-08T08:08:39.826Z)

```
1      ---
2      name: pr-531-single-pass-stitch
3      description: PR
4      metadata:
5        node_type: memory
6        type: project
7        originSessionId: 46eabfd2-0e71-4ff5-a68d-c76f474d4ff3
8      ---
9
10     **PR [#531](https://github.com/Khelsutra/badminton-highlight-indexer/pull/531)** ?
11     `perf(stitch): single-pass concat-filter stitch for a cheaper watermark (#521)`. Branch
12     `perf/521-single-pass-stitch` off master; the deferred Part-3 of [[pr-525-compile-on-l4]].
13
14     **What it does:** opt-in `stitching.single_pass_stitch` (default **OFF**). When on,
15     `VideoCompiler._compile_single_pass` stitches the reel in ONE ffmpeg pass ? precise `-ss/-t`
16     inputs ? `concat` **filter** ? watermark drawn once ? a single encode ? instead of per-clip re-
17     encode + concat-copy. Falls back to the per-clip path on any failure; shared
18     `_resolve_padded_window` validates segments up front; `_SINGLE_PASS_MAX_SEGMENTS` cap guards the
19     N-input graph. Rides the compile-spec to the L4 like `encoder`.
20
21     **Key finding (why the brief's framing was corrected):** a burned-in watermark re-
22     encodes EVERY frame ? that encode is an irreducible floor (a no-watermark stream-copy stitch is
23     ~1.1 s vs ~73 s with it; ~98?99 % of stitch cost IS the watermark re-encode). So neither pre-
24     render (Option A) nor stream-copy-then-concat-re-encode (Option B) removes the floor; the win is
25     per-clip OVERHEAD. Measured on a real 4K 10-bit HEVC clip (RTX 2070S): single-pass (Option C)
26     **?14 % libx264 / ?11 % NVENC, frame-accurate**, smaller output. Option B was rejected (stream-
27     copy snaps to keyframes ? reel 48.1 s vs 40.0 s, a content regression); Option A helps only
28     libx264 (+4 % on NVENC).
29
30     **Verified:** full suite 1889 passed / coverage 85.87 %; ruff+mypy clean; real-ffmpeg
31     e2e (frame-accurate, 4K+audio, watermark rendered). **Owner-gated:** confirming run on the L4 +
32     golden corpus (machine-side split). Doc: `docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md`
33     §3.1.
34
35     **Process note (2026-07-08):** built in an isolated worktree; a concurrent codex session
36     HIJACKED the first worktree (checked out `codex/issue-523-sd-reel-review` + reset it), wiping my
37     UNCOMMITTED changes. Recovered from conversation history, re-applied onto the then-current
38     master (which had gained #527/#528), and **committed immediately**. Lesson: in this high-
39     concurrency checkout, commit early ? uncommitted worktree edits are not safe. Related: my
40     flagged NVENC-probe bug became [[pr-527-nvenc-probe-resolution]] (merged before this PR).
```

290. assistant (2026-07-08T08:08:49.347Z)

The memory system normalized the frontmatter (description got truncated to "PR"); the body is

intact. Let me fix the description and update the body for default-on:

291. user (2026-07-08T08:08:52.503Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\pr-531-single-pass-stitch.md has been updated successfully. (file state is current in your context ? no need to Read it back)

292. user (2026-07-08T08:08:58.095Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\pr-531-single-pass-stitch.md has been updated successfully. (file state is current in your context ? no need to Read it back)

293. user (2026-07-08T08:09:07.998Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\pr-531-single-pass-stitch.md has been updated successfully. (file state is current in your context ? no need to Read it back)

294. assistant (2026-07-08T08:09:14.258Z)

Now the MEMORY.md index line for #531. Let me read and update it:

295. user (2026-07-08T08:09:15.070Z)

```
1      # Memory index
2
3      - [PR #525 ? compile-on-L4 (#521)](pr-525-compile-on-l4.md) ? config-driven
phase?machine placement + L4 NVENC stitch; Part 3 (cheaper watermark) deferred.
4      - [PR #527 ? NVENC probe 64x64?256x256](pr-527-nvenc-probe-resolution.md) ? sub-minimum
probe silently defeated #521's NVENC-on-L4 win; MERGED to master (65ee4c6); STILL awaits owner
live-GPU confirm; monitoring follow-up #529.
5      - [PR #528 ? worker mirror-required](pr-528-worker-mirror-required.md) ? #521 follow-up:
L4-worker stitch now fails on a lost reel mirror-upload (was silent 404-on-success); found by
adversarial review; **MERGED to master** (008a5d4).
6      - [PR #531 ? single-pass stitch (#521 Part-3)](pr-531-single-pass-stitch.md) ? cheaper
watermark via opt-in single-pass concat-filter stitch (default OFF); measured ?11??14% both
encoders, frame-accurate; OPEN, awaits owner L4/corpus confirm.
7      - **Release baseline (2026-07-08):** master `65ee4c6` carries #521 (compile-on-L4) + its
two follow-up fixes #528 (mirror-required) + #527 (NVENC 256px probe). Owner cut/cutting a
release folding all three. Real L4/NVENC exec still unverified (needs a worker image rebuilt
from this master; deployed `:latest` predated the compile subcommand).
8
```

296. user (2026-07-08T08:09:21.173Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

297. assistant (2026-07-08T08:09:31.084Z)

Memory updated. Let me do a final CI check on the updated PR, then summarize.

298. user (2026-07-08T08:09:34.568Z)

```
pass    Built wheel smoke (install + launch)
pass    Dependency license scan (no copyleft in runtime deps)
pass    Golden fixtures cross-OS (macos-latest)
pass    Golden fixtures cross-OS (ubuntu-latest)
pass    Golden fixtures cross-OS (windows-latest)
pass    Leak scan (PII / attribution guard)
```

```
pass    Markdown link check (live docs)
pass    Mypy (lenient baseline)
pass    Ruff lint
pass    Syntax sweep + import smoke (no GPU stack)
pending Full test suite (offline, mocked)
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

299. assistant (2026-07-08T08:09:56.599Z)

Done. `single\_pass\_stitch` is now **default ON**, pushed to PR #531 as a distinct, revertable second commit.

****One correction you'll want to know:**** GitHub's API says ****PR #531 is still OPEN, not merged**** (`mergedAt: none`, and `single\_pass\_stitch` isn't on `origin/master` even after a hard fetch). So "merged" didn't match the actual state. Because it's still open, I put the default-flip ****into #531 itself**** rather than spin up a second PR ? that keeps it "one PR referencing #521" and means ****merging #531 ships single-pass as the default****. If you actually merged something else or want the flip as its own PR, say so and I'll re-cut it.

****What changed (commit `b98209f`):****

- `StitchingConfig.single\_pass\_stitch: False ? True` ? single-pass is now the product default, so the next release picks it up.
- The per-clip path remains the ****automatic fallback**** on any single-pass failure (ffmpeg error, timeout, segment cap) ? that's what makes default-on safe.
- The `VideoCompiler` ***constructor*** arg stays `False` (the conservative primitive default; orchestration passes the config value explicitly) ? this keeps the ~10 unit tests that construct `VideoCompiler` directly to exercise the per-clip/fallback path valid, no churn.
- One test genuinely needed updating: a padding regression test asserted the per-clip `-ss/-to` form; it now asserts the single-pass default's `-ss 8.0 -t 9.0` input-seek form for the same `[8.0, 17.0]` window.
- Doc `PHASE\_PLACEMENT.md` §3.1 updated (status default-ON; fallback keeps it safe; L4/corpus confirmation stays owner-gated and matters a bit more now it's the default).

****Verified (default-flip re-run):****

- Full suite ****1890 passed, 0 failed****, coverage ****85.87%**** (?70) . ruff + mypy clean.
- CI on the updated PR: ****all checks green**** except the full mocked suite still running (reproduced locally).

****Residual (unchanged, honest):**** the confirming run on the real ****L4 + golden corpus**** stays owner-gated. Since single-pass is now the ***default***, that confirmation matters slightly more ? but if it ever misbehaves on real hardware, the per-clip fallback still produces a reel, so it degrades rather than breaks.

PR: <https://github.com/Khelsutra/badminton-highlight-indexer/pull/531>

300. user (2026-07-08T08:26:24.335Z)

merged

301. assistant (2026-07-08T08:26:37.076Z)

Let me verify authoritatively this time ? the last two "merged" signals didn't match GitHub's state, so I'll check the API and confirm it actually landed on master.

302. user (2026-07-08T08:26:43.733Z)

===PR #531 state (GitHub API)===

state: MERGED

mergedAt: 2026-07-08T08:17:45Z

mergeCommit: 5880c26bb0599ef37a3251209a4bd7764bd91be4

origin/master: 5880c26 perf(stitch): single-pass concat-filter stitch for a cheaper watermark (#521) (#531)

===single_pass_stitch default on master now?===

386: single_pass_stitch: bool = True
Shell cwd was reset to C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-indexer

303. assistant (2026-07-08T08:27:06.378Z)

Confirmed ? this time it genuinely merged: ****PR #531 is MERGED**** (`mergedAt 2026-07-08T08:17:45Z`, merge commit `5880c26`), `origin/master` now carries it, and ****`single\_pass\_stitch: bool = True`**** is on master. The default-on landed; the next release will pick it up.

Let me update the memory record to MERGED and clean up my worktree.

304. user (2026-07-08T08:27:11.355Z)

The file C:\Users\avidu\claude\projects\C--Users-avidu-Projects-khehsutra-guru-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

305. user (2026-07-08T08:27:20.230Z)

The file C:\Users\avidu\claude\projects\C--Users-avidu-Projects-khehsutra-guru-badminton-highlight-indexer\memory\pr-531-single-pass-stitch.md has been updated successfully. (file state is current in your context ? no need to Read it back)

306. assistant (2026-07-08T08:27:41.378Z)

Now let me clean up my merged worktree and branch (good hygiene in this heavily-worktree'd repo):

307. user (2026-07-08T08:27:45.914Z)

===remove my merged worktree===
removed (or already gone)
===delete merged local branch===
Deleted branch perf/521-single-pass-stitch (was b98209f).
===confirm my worktree/branch gone===
worktree gone ?
branch gone ?

308. assistant (2026-07-08T08:27:57.914Z)

All wrapped up. ?

****Merged and landed:**** PR #531 is in master (`5880c26`), and ****`single\_pass\_stitch = True`** is now the default ? so the next release ships the single-pass concat-filter stitch by default (~11?14% faster, frame-accurate, both encoders), with the per-clip path as the automatic fallback if it ever fails.

Cleaned up my merged worktree and local branch. Memory record updated to MERGED.

****One honest residual carried forward**** (tracked in the doc §3.1 and project memory): the confirming run on the real ****L4 + golden corpus**** is still owner-gated. It matters a little more now that single-pass is the default ? but the fallback means a failure on real hardware degrades to the old per-clip stitch rather than breaking. That's the one thing worth closing when you next have the GPU box + corpus in hand.