

claude-windows / 75d3d121-f240-4293-afa5-b436a379ad31

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\75d3d121-f240-4293-afa5-b436a379ad31.jsonl`  
- SHA-256: `7da25009fd6bc64c958d463207f5bbdc0f7e30e0252330508aae3645ea022485`  
- Source modified: `2026-07-08T06:52:44+00:00`  
- Imported at: `2026-07-08T15:59:19+00:00`  
- project: `C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer`  
- session_id: `75d3d121-f240-4293-afa5-b436a379ad31`
```

Transcript

1. user (2026-07-08T05:04:37.601Z)

Implement GitHub issue #521 (Khelsutra/badminton-highlight-indexer): "perf/scaling: run compile/stitch on the L4 worker; make phase?machine placement config-driven". Read issue #521 and #520 first (`gh issue view 521 -R Khelsutra/badminton-highlight-indexer``), plus the evidence in `docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md`` §9.

Repo / environment

```
- Real code lives HERE: `C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer`  
(the sibling `C:\Users\avidu\Projects\badminton-highlight-indexer` is an empty STUB ? ignore it).  
- Windows + PowerShell primary; `gcloud`?googleapis egress works but general web egress is blocked (curl to non-Google hosts returns HTTP 000). `gh` and `git` to GitHub DO work. Authed as avidullu, project `gen-lang-client-0306026826`, region `asia-southeast1`.
```

Why (measured, 2026-07-08 live CUJ on rev rally-control-plane-00019)

Compile/stitch runs IN-PROCESS on the CPU-only 2-vCPU control-plane and took `**~12m55s**` for an 11-clip / ~80s reel from a 1.83 GiB 4K source ? `**~45?90s per ~7s clip**`, because the per-clip branding watermark forces a full 4K re-encode with no NVENC. Validation (183s) ALSO runs on the control-plane; both are serialized by a mutex, so one run pins the control-plane for ~16 min of heavy CPU ? won't scale past ~5 concurrent users. Detection already dispatches to an L4 Cloud Run Job and is fast.

Current architecture (files to study first)

```
- `backend/orchestration.py`:  
  - `_execute_compile(cfg, db, job, progress)` ? the compile/stitch handler (kind='compile' job): resolves segments, runs `backend.main.VideoCompiler` (ffmpeg trim+watermark+concat) under `_LOCAL_COMPUTE_LANE`, mirrors the reel to `reel_object_uri()` = `/highlights/<video_id>/<name>`.  
  - `_resolve_compile(...)` ? already documented as "run at SUBMIT AND in the WORKER (re-resolve)", so it's built to run worker-side.  
  - `_maybe_dispatch_remote(...)` ? the DETECT dispatch path (control_plane ? Cloud Run Job), emits `--skip-validation`, forwards `indexing.wasb.*`, bucket-polls a `done` marker. THIS is the pattern to mirror for compile.  
  - `_cloud_available(cfg)`, `reel_object_uri`, `signed_reel_url`.  
- `backend/worker/__main__.py` ? the worker entrypoint (`infer` subcommand today). Add a `compile` subcommand (or a `--phase=compile` mode) that runs `_execute_compile` worker-side and writes the reel + a done-marker.  
- `backend/compute/cloud_run.py` ? Cloud Run Job dispatch + bucket-poll plumbing.  
- `backend/config/models.py` ? config schema (`deployment.compute_target`, `indexing.wasb.*`, `StitchingConfig`); add the phase-placement config here.  
- The cloud-serving preset (grep `cloud-serving` under `backend/config/`) ? where `compute_target=cloud_run` is set.  
- `backend/main.py::VideoCompiler` (the ffmpeg wrapper; watermark path).  
- Tests: `tests/test_api_cloud_dispatch.py`, `tests/test_worker.py`,
```

`tests/test\_compute\_backend.py`, `tests/test\_async\_jobs.py` ? mirror these for the compile-dispatch path.

Deliver (3 parts)

1. ****Dispatch compile/stitch to the L4 worker.**** Add a `kind='compile'` Cloud Run dispatch path mirroring `\_maybe\_dispatch\_remote`: the control-plane submits a compile job to the `rally-worker` Job with the video_id + segment_ids + output_filename + locale, the worker runs `\_execute\_compile` (re-resolving via `\_resolve\_compile`), writes the reel to `/highlights/<video\_id>/<name>` + a `.done` marker, and the control-plane bucket-polls + returns the same `{status, filepath, download\_url}` shape the SPA already expects. Prefer NVENC (`h264\_nvenc`/`hevc\_nvenc`) on the L4 for the trim/watermark re-encode; keep libx264 as the CPU fallback.
2. ****Config-driven phase?machine placement.**** Add a per-phase placement knob (e.g. `deployment.phase\_placement: {validation: control\_plane, segment: cloud\_run\_l4, compile: cloud\_run\_l4}` with values `control\_plane | cloud\_run\_l4`), defaulting to today's behavior (validation+compile on control-plane, segment on cloud_run when compute_target=cloud_run) so nothing changes unless configured. Wire compile + detect to consult it; set compile?cloud_run_l4 in the cloud-serving preset. Goal: move a phase between machines by config edit, no code change.
3. ****Cheaper watermark**** (the dominant term): pre-rendered/overlay-filter or an NVENC filter path so the stitch isn't a full CPU re-encode. If large, land 1+2 first and file/track 3 separately.

Acceptance

- Compile can be dispatched to the L4 by config; control-plane stitch remains available as a config option (both paths tested, mocked ? mirror `tests/test\_api\_cloud\_dispatch.py`).
- Flipping `phase\_placement` moves validation / segment / compile between machines with no code change.
- New + existing tests green; `ruff check .` and `mypy` clean (match CI); `python -m tools.leak\_scan` and `python -m tools.check\_md\_links` green; run the FULL suite (e.g. `python run\_all\_tests.py` / `pytest`), not just changed files; keep coverage ? the repo's `--cov-fail-under` floor.
- Update `docs/COMPUTE\_DECOUPLED\_SERVING/` docs for the new config + phase-placement model.

Branching safety (concurrent sessions share this checkout ? follow every time)

- The shared checkout may have active sibling WIP (e.g. `feat/validation-result-cache`) and master can move under you. Work in an ISOLATED clone/worktree; never touch the shared working tree or sibling branches.
- `git fetch origin` then branch EXPLICITLY from the remote: `git checkout -b feat/521-compile-on-l4 origin/master`. Re-verify right before commit: `git branch --show-current` + `git log --oneline origin/master..HEAD` must list ONLY your commits (re-cut from `origin/master` if a sibling commit appears).
- Stage explicit paths (`git add <files>`), never `git add -A`/`.`/`-u`.
- Open a PR against `master`, reference #521 (and note it's driven by the 2026-07-08 live-CUJ evidence in the §9 report). End the PR body with the Claude Code attribution line.

Deploy is out of scope ? implement + test + PR only. Ask before any `gcloud run deploy` / traffic change.

2. assistant (2026-07-08T05:04:56.568Z)

I'll start by following the branching-safety and repo-freshness conventions: gather context (issues, evidence, memory/handoff) and check the repo state before making any changes. Let me kick off several independent reads in parallel.

3. user (2026-07-08T05:05:18.796Z)

title: perf/scaling: run compile/stitch on the L4 worker; make phase?machine placement config-driven
state: OPEN
author: avidullu (Avi)
labels: bug, enhancement
comments: 0
assignees:

projects:
milestone:
number: 521
--

Problem

The `compile/stitch` phase runs `in-process` on the CPU-only 2-vCPU control-plane
(`backend/orchestration.py:_execute_compile` ? `_LOCAL_COMPUTE_LANE` ffmpeg`). `Validation` also runs on the control-plane. On the 2026-07-08 live CUJ (rev `rally-control-plane-00019``, clip `"First CUJ.MP4" ? 1.83 GiB / 174.17 s / 4K-class`):

```
| Phase | Machine | Wall |  
|---|---|---|  
| Validation (#512 shared-decode) | control-plane 2 vCPU | **183 s** |  
| Detection (WASB) | L4 worker | 161.6 s |  
| **Compile / stitch** (11 clips ? ~80 s reel) | **control-plane 2 vCPU** | **~12m55s** |
```

Stitch was `~45?90 s per ~7 s clip` because the `per-clip branding watermark` forces a full 4K re-encode with no NVENC on 2 vCPU. So a single run pins the control-plane for `~16 min` of heavy CPU (validation + stitch), serialized by `_LOCAL_COMPUTE_LANE``. At even a handful of concurrent users the control-plane saturates and the API becomes unusable ? estimated `>~5` concurrent users renders it useless.

Ask

- `Move compile/stitch to the L4 worker` (hardware `NVENC` + 8 vCPU), the way detection already dispatches. The worker already fetches the source, and `_execute_compile`` is designed to re-resolve in the worker (`_resolve_compile`` is explicitly `"run at SUBMIT **AND** in the WORKER"`). Add a `kind='compile`` Cloud Run dispatch path.
- `Make phase?machine placement configurable` ? a per-phase target knob (`validation``, `segment/detect``, `compile/stitch`` ? `control_plane | cloud_run_l4 | <other>``) so phases can be relocated `by config, no code change`. Mirror the existing `deployment.compute_target`` + `_maybe_dispatch_remote`` forwarding pattern (`COMPUTE_DECOUPLED_SERVING`).
- `Cheaper watermark` ? pre-rendered overlay / NVENC filter path so stitch isn't a full CPU re-encode (the dominant term above).

Acceptance

- Compile can be dispatched to the L4 by config; control-plane stitch stays available as a config option.
- Flipping config moves validation / detect / compile between machines `without a code change`.
- Measured stitch wall for this reel drops from `~13 min (2 vCPU)` to sub-minute-class on L4.

Evidence

`docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md`` §9 (2026-07-08 live CUJ). Compile timeline: `POST /api/compile` 04:37:26 ? 11x per-clip 4K re-encode+watermark (~67 s each) ? concat ~2 s ? reel mirrored 04:50:21 (`first-cuj-ks-highlight-reel.mp4`, 581.8 MB).`

`_Filed from the 2026-07-08 live-CUJ phase breakdown._`

4. user (2026-07-08T05:05:20.540Z)

title: Observability: step-level debug tracing (server + UI) to a low-latency sink
(bucket/service)
state: OPEN
author: avidullu (Avi)
labels: enhancement
comments: 0
assignees:
projects:
milestone:
number: 520
--

Problem

Observing a run today means scraping Cloud Run stdout **after the fact** via ``gcloud logging read``, then hand-correlating timestamps across **three** separate streams ? the control-plane service, the worker Job, and the in-process compiler ? and there is **no UI-side signal at all**. Reconstructing the phase timeline for the 2026-07-08 live CUJ (upload ? validate ? dispatch ? detect ? ingest ? compile) required manually diffing log timestamps across those streams.

Ask

A **debug/trace mode** (config- or env-gated, **off by default**) where every pipeline step ? and key UI actions ? emits a **structured event** to a **low-latency sink**, so a run can be observed live and replayed:

- **Server phases:** upload (per-part / complete), validation (per-validator wall + verdict), dispatch (job/exec id + cold-start), worker (pull, skip-validation, WASB phase), ingest, compile (per-segment trim, concat, mirror-upload).
- **UI events:** upload start/complete, Process click, poll states, Compile click, download.
- **Event shape:** `{ts, run_id, video_id, phase, step, machine, wall_ms, meta}`.
- **Sink (configurable):**
 - GCS ? one JSONL per run at ``gs://<output_root>/_debug/<video_id>/trace.jsonl`` (cheap, durable, replayable), and/or
 - a lightweight collector / Pub/Sub topic for **live** viewing (e.g. an ``/api/debug/trace/<video_id>`` SSE stream, or push to an existing observability service).
 - Gate behind e.g. ``observability.debug_trace: true`` (+ sink target); never on by default (PII / volume).

Why now

The live CUJ (``docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md`` §9) had to be timed **by hand** across control-plane + worker + compiler logs; a first-class per-run trace (server **and** UI) would make the whole phase breakdown a single artifact and enable low-latency live observation.

`_`Filed from the 2026-07-08 live-CUJ observability gap.`_`

5. user (2026-07-08T05:05:23.491Z)

From <https://github.com/Khelsutra/badminton-highlight-indexer>

```
* [new branch]      docs/validation-latency-live-cuj-2026-07-08 -> origin/docs/validation-latency-live-cuj-2026-07-08
```

```
---STATUS---
```

```
On branch master
```

```
Your branch is up to date with 'origin/master'.
```

```
Untracked files:
```

```
  (use "git add <file>..." to include in what will be committed)
```

```
    score_golden_gate.py
```

```
    score_local.py
```

```
nothing added to commit but untracked files present (use "git add" to track)
```

```
---BRANCH---
```

```
master
```

```
---LOG---
```

```
699d85b feat(validation): cache validation verdict by video_id to skip re-validation on reupload (#518)
```

```
71c88ef docs(runlogs): validation-latency closure report for #512 + #513 (cloud-serving) (#519)
```

```
115b83d docs(runlogs): record live smoke of the crash-before-marker fast-fail (#515) (#517)
```

```
ef41473 fix(cloud-run): fail fast when a worker execution terminates without a done-marker (#515)
```

```
f5c9dd9 chore(ui): re-vendor SPA bundle with the hosted compute-picker cloud default (#516)
```

6. assistant (2026-07-08T05:05:27.779Z)

Now let me read the evidence report §9, check memory/handoff, and locate the key architecture files.

7. user (2026-07-08T05:05:31.992Z)

```
===MEMORY DIR===
total 4
drwxr-xr-x 1 avidu 197609 0 Jul  8 10:34 .
drwxr-xr-x 1 avidu 197609 0 Jul  8 10:34 ..
===REPORT===
DEPLOY_REPORT_cloud-serving_2026-07-04.md
DEPLOY_REPORT_cloud-serving_a1-cuj_2026-07-04.md
DEPLOY_REPORT_cloud-serving_a1-e2e_2026-07-04.md
DEPLOY_REPORT_cloud-serving_alpha-alliv1_2026-07-06.md
DEPLOY_REPORT_cloud-serving_cp-rebuild_2026-07-05.md
DEPLOY_REPORT_cloud-serving_gpu-resident-nvdec_2026-07-07.md
DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md
README.md
RUNLOG_cloud-serving_exec-status-fastfail_2026-07-07.md
RUNLOG_truly-local_TEMPLATE.md
```

8. user (2026-07-08T05:05:33.943Z)

```
1      # DEPLOY REPORT ? cloud-serving, validation-latency closure (#512 shared-decode + #513
worker skip-revalidation), 2026-07-07
2
3      > Durable **closure record** for the two validation-latency fixes now live on the
`cloud-serving` path: **#512** (`perf(validators)`: decode `scene_cut`/`blur`/`relevance` in ONE
shared pass instead of ~125 random per-sample seeks across 3 captures) and **#513**
(`fix(serving)`: the worker skips its redundant re-validation when the control-plane already
validated + dispatched, via `--skip-validation`). Both shipped in the live image `rally-
worker@sha256:8045f54e?` (built from `origin/master` @ `f5c9dd9`; #512 `20862d7`, #513 `23c1868`
are ancestors). Measures the before/after on 3 clips spread across codecs/resolutions, all
against the real serving infra. Project `gen-lang-client-0306026826`, region `asia-southeast1`.
Honest-limits (incl. the owner-gated live CUJ) in $6.
4
5      ## 1. TL;DR
6
7      | Stage | Result |
8      |---|---|
9      | Live image carries #512+#513 | ? control-plane rev `00018-qvz` + worker Job both on
`sha256:8045f54e?` (= `f5c9dd9`, ancestors `20862d7`/`23c1868` verified) |
10     | Baseline captured (pre-fix) | ? from prod logs on rev `00017-qbd` (image `09d2edf7` =
`66756ad`, pre-#512/#513) ? no re-run needed |
11     | **#513 worker skip-revalidation** | ? real dispatch-style exec: Video 1 validation
**SKIPPED (0 s)**, runs to **SUCCESS / 13 rallies** ? the blur-reject that failed the live CUJ
is **gone** |
12     | **#512 shared-decode speedup (real L4)** | ? Video 1 re-validation **169.3 s ? 69.7 s
(2.43x)** on identical worker hardware, decision byte-identical (sharpness 11.9 < 20.0 both) |
13     | **#512 speedup (off-box A/B, breadth)** | ? 1080p H.264 **1.95x**, 4K HEVC **3.10x**,
5.3K 10-bit HEVC **3.23x** ; every validator decision unchanged OLD?NEW |
14     | Live control-plane "after" number | ? owner-gated (Cloudflare Access) ? projected
**-412 s ? ~170 s (2?3 min)** ; recipe in $6 |
15
16     **Bottom line:** the two fixes do exactly what they claim on live serving hardware.
**#513** turns the soft 5.3K clip from a **hard FAIL** (worker re-validated at its default
`min_blur=20`, rejected sharpness 11.9, `rc2`, no reel) into a **PASS with 13 rallies** by
trusting the control-plane's gate. **#512** cuts the frame-sampling validators ~**2?3x** (2.43x
measured on the L4; 1.95?3.23x off-box across codecs) with **zero decision drift** OLD?NEW. The
only number still owner-pending is the live control-plane validation wall on rev 00018 ? its
components are measured and it projects to ~2?3 min from the ~6m52s baseline.
17
18     ## 2. Resources
19     - **Control-plane:** `rally-control-plane` Cloud Run **service**, rev **`rally-control-
plane-00018-qvz`** (100% traffic), image `asia-southeast1-docker.pkg.dev/gen-lang-
client-0306026826/rally/rally-worker@sha256:8045f54ea9fcbf?` (built from `f5c9dd9`),
`DEPLOYMENT_PROFILE=cloud-serving`, CPU-only **2 vCPU**, edge-auth ON (`edge_auth_enabled=true`,
`cf_team_domain=khelsutra.cloudflareaccess.com`).
```

```

20 - **Worker:** `rally-worker` Cloud Run **Job**, same image `sha256:8045f54e?`, 8 vCPU /
32 GiB / 1x NVIDIA **L4**, `WASB_WEIGHTS=/gcs/models/wasb_badminton_best.pth.tar`,
`WASB_DECODE_BACKEND=nvdec_resident`, GCSFuse mount of `gs://khelsutra-rally-serving`.
21 - **Baseline (pre-fix) surface:** control-plane rev **`00017-qbd`** / worker image
`sha256:09d2edf7?` (= `66756ad` = #507+#508+#509, both **pre-#512/#513**)? the source of the
prod baseline logs below (2026-07-07, before the 16:15 rev-00018 cutover).
22 - **Corpus:** Video 1 `gs://khelsutra-rally-
serving/videos/uploads/702636eed76145e88aa59c3000c7e415/Video 1 - Badminton.MP4` (1.41 GiB,
5312x2988 yuv420p10le HEVC, 3019 frames, MD5 `8fe6b719603dd7a7b2b6ef543139a4a3`);
`gs://khelsutra-rally-corpus/videos/golden_1080p/Boxhill_Doubles.mp4` (1080p H.264);
`gs://khelsutra-rally-corpus/videos/golden/GX010128.mp4` (2.18 GiB, 4K HEVC, 45298 frames).
23
24 ## 3. Exact commands
25 ``bash
26 # --- Provenance: prove the live image carries both fixes ---
27 gcloud run services describe rally-control-plane --region asia-southeast1 \
28 --format="value(status.latestRevisionName,
spec.template.spec.containers[0].image)" # 00018-qvz ?@sha256:8045f54e?
29 git merge-base --is-ancestor 20862d7 f5c9dd9 && echo "#512 in live image" # yes
30 git merge-base --is-ancestor 23c1868 f5c9dd9 && echo "#513 in live image" # yes
31
32 # --- #513 proof: worker A/B on the LIVE image (Video 1 already staged in the serving
bucket) ---
33 V1="gs://khelsutra-rally-serving/videos/uploads/702636eed76145e88aa59c3000c7e415/Video 1
- Badminton.MP4"
34 # Run A ? WITH --skip-validation (exactly what orchestration._maybe_dispatch_remote
emits on dispatch):
35 gcloud run jobs execute rally-worker --region asia-southeast1 --async \
36 --args="^@@^infer@@--video-uri=$V1@@--out-uri=gs://khelsutra-rally-
serving/_valclose/v1-skipval@@--segmenter=wasb_hybrid@@--skip-validation@@--
set=deployment.compute_target=local_gpu@@--set=indexing.detector_impl=native@@--
set=indexing.skip_ai_handoff=true@@--set=indexing.wasb.batch_size=64@@--
set=indexing.wasb.prefetch_batches=4@@--set=indexing.wasb.timeout_sec=0"
37 # Run B ? WITHOUT --skip-validation (worker re-validates with the NEW #512 validators ?
still blur-rejects):
38 gcloud run jobs execute rally-worker --region asia-southeast1 --async \
39 --args="^@@^infer@@--video-uri=$V1@@--out-uri=gs://khelsutra-rally-serving/_valclose/v
1-reval@@--segmenter=wasb_hybrid@@--set=deployment.compute_target=local_gpu@@--
set=indexing.detector_impl=native@@--set=indexing.skip_ai_handoff=true"
40
41 # --- #512 timing: bracket "pulled ? -> in.mp4" ? "validation FAILED/SKIPPED" ?
"Phase 2 (WASB) completed" ---
42 gcloud logging read 'resource.type="cloud_run_job" AND
labels."run.googleapis.com/execution_name"="rally-worker-<exec>" \
43 --freshness 1d --format="value(timestamp, textPayload)"
44
45 # --- #512 off-box A/B (isolated clones; run_all decodes the SAME sample indices in both
trees) ---
46 git clone --local --no-hardlinks <repo> repo_old && git -C repo_old checkout 23c1868 #
pre-#512 (per-sample seeks)
47 git -C repo_old worktree add --detach repo_new f5c9dd9
# live image src (shared pass)
48 python valbench.py <repo_old|repo_new> <video> <label> 2 # times
registry.run_all(video, load_config())
49 ``
50
51 ## 4. Run results
52
53 ### 4a. Video 1 (5312x2988 10-bit HEVC, 3019 frames) ? the clip that motivated #512/#513
54
55 | Metric | BEFORE ? pre-fix (rev `00017`, img `09d2edf7`) | AFTER ? #512+#513 live (rev
`00018`, img `8045f54e`) |
56 |---|---|---|
57 | Control-plane validation | **411?412 s** (6m52s) ? passed (baked `min_blur=8.0`) | ?
live owner-gated; **projected ~170 s (2?3 min)** |

```

```

58 | Worker (re-)validation | **169.3 s** ? ? **BLUR-REJECT** (sharpness 11.9 < default
20.0) | **0 s ? SKIPPED** (--skip-validation`, #513) |
59 | &nbsp;&nbsp;&nbsp;? counterfactual: worker re-val if NOT skipped | *(169.3 s, OLD
validators)* | **69.7 s** (NEW #512 validators, **2.43***, same 11.9<20 decision) |
60 | End-to-end | ~10m31s wall ? **FAILED** (0 rallies, no reel) | worker ? **SUCCESS** in
~2m38s exec ? **13 rallies** (reel = owner CUJ) |
61 | Rallies produced | **0** (rejected before detection) | **13** (`report.json`
`status=success`, 13 segments) |
62 | Pass / fail | ? **FAIL** | ? **PASS** |
63
64 **Prod baseline evidence (rev 00017, 2026-07-07, unmodified):** control-plane `13:49:18
Running validations ? 13:56:10 Video passed checks` = 411.4 s (and `14:03:56 ? 14:10:48` = 411.9
s). Worker exec `rally-worker-2lbz2` (dispatched by the 14:03 run, OLD image): `pulled 14:11:35
? validation FAILED: blur_check ? sharpness 11.9 below 20.0 @14:14:24 ? exit(2)` ? a
**successful GPU dispatch that failed after the fact**, the exact CUJ #513 fixes.
65
66 **After evidence (live image `8045f54e`, 2026-07-07):**
67 - `rally-worker-k29rk` (**Run A, --skip-validation**): `pulled 17:09:24 ? [worker]
validation SKIPPED (--skip-validation) 17:09:27 ? native WASB (torch 2.6.0+cu124, cuda True) ?
Phase 2 (WASB) completed in 100.9 s, 13 candidate windows (1510/3019 visible) ? emitted 13
rallies ? exit(0)`. Exec wall 157.7 s; `gpu_active`?100.9 s (WASB);
`bytes_out`=`report.json`+`intervals.csv`. Output
`gs://?/_valclose/v1-skipval/outputs/8fe6b719?/`, `status=success`, 13 segments.
68 - `rally-worker-t8llx` (**Run B, re-validation on the NEW image**): `pulled 17:09:29 ?
validation FAILED: blur_check ? sharpness 11.9 below 20.0 @17:10:39 ? exit(2)`. Re-val wall
**69.7 s** vs the OLD image's 169.3 s = **2.43x on identical L4 hardware**, and the **decision
is byte-identical** (11.9 < 20.0) ? #512 is pure speed, zero behavior change. (This run also
shows *why* #513 is needed: absent the skip, the worker's default `min_blur=20` still rejects
the clip the control-plane passed at 8.0.)
69
70 ### 4b. Off-box #512 A/B ? codec/resolution breadth (16-core local box, cv2 4.13.0,
warm, `run_all`)
71
72 Isolated clones at `23c1868` (OLD, per-sample seeks) vs `f5c9dd9` (NEW, shared pass);
identical sample indices ? a fair A/B. **Every validator decision is identical OLD?NEW** on
every clip (no PASS/FAIL drift from the fix).
73
74 | Clip | codec / res | OLD (warm) | NEW (warm) | Speedup | decisions OLD?NEW |
75 |---|---|---|---|---|---|
76 | Boxhill_Doubles | 1080p H.264 | 10.91 s | 5.60 s | **1.95x** | 6/6 identical (all
PASS) |
77 | GX010128 | 4K HEVC (45298 fr) | 35.76 s | 11.54 s | **3.10x** | 6/6 identical (all
PASS) |
78 | Video 1 | 5.3K 10-bit HEVC | 173.20 s | 53.62 s | **3.23x** | 6/6 identical (all
PASS?) |
79
80 ? On this Windows/cv2 box the 10-bit-HEVC decode yields a Laplacian variance **above**
20 so blur PASSES off-box, whereas the Linux serving stack measures **11.9** and rejects. That
platform decode-stack difference is **orthogonal to #512** (which changes only *which frames are
sampled*, identically in both trees); the authoritative blur verdict and the real speedup come
from the L4/prod ($4a). Off-box contributes only the OLD?NEW **ratio**, which is decode-stack-
independent.
81
82 The win grows with per-frame decode cost (random-seek re-decode is what #512 removes),
and it **triangulates** with the real-L4 Video 1 A/B (2.43x): the off-box ratios bound the
control-plane projection.
83
84 ### 4c. Control-plane projection (2-vCPU)
85 Baseline **~412 s**. Applying the L4-measured **2.43x** ? **~170 s (~2m50s)**; the off-
box ratios (1.95?3.23x) bracket **~2?3 min**. Matches the pre-fix design estimate (~6.5 min ? ~2
min). Exact live number is owner-gated ($6).
86
87 ## 5. Rollout state
88 - **Already deployed ? this is a measurement/closure record, not a deploy.** #512+#513
went live with the `8045f54e` image cutover (control-plane rev `00018-qvz`, worker Job same

```

digest) on 2026-07-07. No infra change was made for this report.

```

89 - **Rollback (pre-existing levers):** control-plane `gcloud run services update-traffic
rally-control-plane --region asia-southeast1 --to-revisions rally-control-plane-00017-qbd=100`;
worker `gcloud run jobs update rally-worker --region asia-southeast1 --image ?/rally-
worker@sha256:09d2edf7?`. Both revert to the pre-#512/#513 image.
90 - **Scratch outputs** written under `gs://khehsutra-rally-serving/_valclose/` (isolated
prefix; safe to delete).
91
92 ## 6. Honest limits / not-verified
93 - **Live end-to-end CUJ is owner-gated.** `POST /api/process` on rev 00018 requires a
Cloudflare-Access JWT (`edge_auth_enabled=true`; a direct call returns **401** `missing Cf-
Access-Jwt-Assertion header`). So the two ? cells ? the **live control-plane validation wall on
rev 00018** and the **stitched reel** ? need the owner to drive one upload (or hand over a CF-
Access JWT / service token). Everything else here is measured on the real serving infra.
**Recipe to close them:** upload "Video 1 - Badminton.MP4" via the SPA (compute=cloud), then
`gcloud logging read '?service_name="rally-control-plane"?("Running validations" OR "passed
checks")'` for the wall, and confirm the dispatched worker exec logs `validation SKIPPED` +
emits rallies + the reel compiles. Expect CP validation ~2?3 min and a PASS.
94 - **#513's worker win is shown via a dispatch-faithful direct exec**, not a control-
plane-initiated dispatch (same auth gate). The `--skip-validation` flag, its emission in
`orchestration._maybe_dispatch_remote`, and its handling in `worker.cmd_infer` are unit-tested
on master; Run A exercises the exact worker code path with the exact flag.
95 - **Rallies, not a reel.** Run A used `skip_ai_handoff=true` (secret-free) so the 13
"rallies" are local candidate windows (no Gemini refine, no stitched mp4) ? this validates
decode+detect+windowing end-to-end and that the clip **passes**; the reel stitch is the owner-
CUJ step above. The 13 segments are its inputs.
96 - **Off-box absolute times are from a Windows 16-core box (cv2/FFmpeg), not the 2-vCPU
Linux control-plane** ? so the **ratios** transfer (the change is algorithmic: seeks?shared
forward-hybrid pass), the absolutes don't, and (per ?) the 10-bit blur *verdict* can differ; the
control-plane absolute is anchored to the prod 412 s baseline, not the local numbers. Video 1
off-box is a cross-check of the authoritative real-L4 2.43x.
97 - **#515** (crash-before-marker fast-fail) is NOT in this image (`ef41473` post-dates
`f5c9dd9`); it's unrelated to validation latency and out of scope for this closure.
98
99 ## 7. Cost
100 - 2 L4 worker execs (Run A 157.7 s + Run B 109.6 s ? 268 s of L4 time) ? **~$0.5?0.7**.
Off-box A/B: **$0** (local). No VM, no image build. GCS read for the 3 downloaded clips (~5.1
GiB) is intra-project.
101
102 ## 8. Closure verdict
103 ? **CLOSED ? #512 and #513 are proven on live cloud-serving hardware.** #513 converts
the motivating soft-focus 5.3K clip from a post-dispatch FAIL into a 13-rally PASS by
eliminating the worker's redundant re-validation (169 s + wrongful blur-reject ? 0 s); #512
speeds the frame-sampling validators **2.43x on the real L4** and **1.95?3.23x off-box across
codecs**, with **zero decision drift** OLD?NEW on every clip. The only residual is the live
control-plane validation wall on rev 00018 (owner-gated, projected ~2?3 min from ~6m52s) ? its
components are all measured. Validation-latency work is done.
104

```

9. assistant (2026-07-08T05:05:41.291Z)

The report ends at §8, but §9 (the 2026-07-08 live CUJ) is referenced. A new remote branch `docs/validation-latency-live-cuj-2026-07-08` just appeared ? §9 likely lives there. Let me check.

10. user (2026-07-08T05:05:43.910Z)

===does §9 exist on the new remote branch?===

9. Update (2026-07-08) ? live SPA CUJ closes the two owner-gated cells (on rev `00019`)

The two ? cells in §4a (the **live control-plane validation wall** and the **end-to-end reel**) were owner-gated on the Cloudflare-Access `POST /api/process`. On **2026-07-08** the owner drove one real SPA upload (compute=**Cloud**) end-to-end, so both are now measured on live serving infra. Mirrors the gpu-resident report's "§5 Update" append ? the body above is unchanged; this

section records what the live run added.

****Two honest deltas from the recipe**** (state moved after this report's 2026-07-07 image cutover):

1. ****Live rev rolled `00018` ? `00019`.** ****Traffic is 100% on `rally-control-plane-00019-96f`**, image `rally-worker@sha256:c6393d27b4ae?` (tag `:latest`), deployed 2026-07-07T17:06Z ? ~51 min after the `00018`/`8045f54e` cutover \$2 documents; the worker Job is on the same `:latest`=`c6393d27` digest (both surfaces match). It's an undocumented forward-bump, but #512 (`20862d7`) and #513 (`23c1868`) are ancestors of `origin/master`, `f5c9dd9` (the `00018` src) is an ancestor of master, and the run ****empirically**** shows both fixes live (one-pass validation wall + worker `validation SKIPPED`). Measured on what is actually serving.

2. ****A fresh clip, not Video 1.**** The upload was ****"First CUJ.MP4"***** ? ****1.83 GiB**** (1,960,043,219 B), source duration ****174.17 s**** (?84 Mbps, 4K-class ? ***heavier*** than Video 1's 1.41 GiB), video_id `2f31b45e?`. So the live wall ****corroborates**** the Video-1 projection on a comparably-heavy clip; it is ****not**** a re-measurement of Video 1's 412 s?~170 s A/B (a different clip has a different absolute wall).

****Full phase breakdown (UTC 2026-07-08, rev `00019`):****

#	Phase	Window	Wall	Machine	Notes
1	Upload (client?server, 80-part multipart)	04:25:56 ? 04:28:51	~2m55s	client uplink	1.83 GiB @ ~84 Mbps ? owner bandwidth, not infra
2	Accept ? validation start (stage)	04:28:52 ? 04:29:22	~30 s	control-plane	stage upload to serving bucket
3	**Validation (#512 shared-decode)**	04:29:22 ? 04:32:25	~3m18s (3m03s)	control-plane 2 vCPU	one-pass scene_cut/blur/relevance; PASSED
4	Dispatch ? worker ready (cold start)	04:32:25 ? 04:33:20	~55 s	Cloud Run Job	image pull + GCSFuse mount + torch/cuda init
5	**Worker skip-validation (#513)**	04:33:20 ? 04:33:24	~4 s	L4	`validation SKIPPED`, 0 re-validation
6	Detection / segment (WASB)	04:33:24 ? 04:36:20	~3m16s	L4 (cuda)	15 candidate windows, 2561 visible pts
7	Ingest + finalize	04:36:20 ? 04:36:24	~4 s	control-plane	intervals?DB, `status=success`
?	(user reviews rallies, clicks Compile)	04:36:24 ? 04:37:26	~62 s idle	?	human-in-the-loop
8	**Compile / stitch**	04:37:26 ? 04:50:21	~12m55s	control-plane 2 vCPU	11 clips, ~67 s/clip 4K re-encode+watermark; concat ~2 s; mirror ~11 s

Server pipeline (process?success, excl. upload+idle): ****7m32s****. Full CUJ (upload?reel): ****~24m25s**** wall.

****Cell 1 ? control-plane validation AFTER (live):**** ****183 s (3m03s)**** on rev `00019`, PASSED ? the #512 shared-decode validators on the real 2-vCPU serving path, in the projected ****2?3 min**** band on a clip ***heavier*** than Video 1 (pre-#512 per-sample seeks were the ~6m52s-class baseline). Evidence: `Running validations 04:29:22.789 ? Video passed checks 04:32:25.780`.

****Cell 2 ? end-to-end reel (live):**** `/api/compile` stitched ****11 of the 15**** detected rallies (0.5 s pad, per-clip branding watermark) from the 174.17 s source ? ****gs://kheutra-rally-serving/highlights/2f31b45e?/first-cuj-ks-highlight-reel.mp4**** (****581.84 MiB****, saved 04:50:10 ****"Dynamic highlight reel saved (581.8 MB)"****, mirrored 04:50:21, served via `/api/highlights` 307). ****Reel ? produced end-to-end on the live control-plane**** ? source fetch ? per-clip trim+watermark ? concat ? mirror-upload, all exercised.

****Bonus ? #513 proof upgraded**** from the \$6 dispatch-faithful ***direct exec*** to a genuine ****control-plane-initiated dispatch**** (`rally-worker-64bhj`): `\_maybe\_dispatch\_remote` emitted `--skip-validation` on a live SPA run and the worker honored it (0 re-validation).

****? Finding ? stitch is the scaling bottleneck (follow-ups filed).**** The compile ran ****~13 min** on the CPU-only 2-vCPU control-plane** (~45?90 s per ~7 s clip: the watermark forces a full 4K re-encode, no NVENC). With validation (183 s) ***also*** on the control-plane and serialized by `\_LOCAL\_COMPUTE\_LANE`, one run pins the control-plane for ~16 min of heavy CPU ? it won't scale past a few concurrent users. Filed: ****#521**** (move compile/stitch to the L4 worker + make phase?machine placement config-driven) and ****#520**** (step-level debug tracing, server+UI, to a

low-latency sink ? this breakdown had to be hand-correlated across 3 log streams).

****Cost (heuristic):**** one L4 detect exec (``rally-worker-64bhj``, ~222 s) ? ****\$0.15?0.55**** (rate-dependent, dominant) + control-plane CPU for validation+stitch (~960 s @ 2 vCPU) ? ****~\$0.05**** + reel egress (581 MB, if downloaded to internet) ? ****~\$0.07**** + storage/ops ****< \$0.02**** + Gemini ****\$0**** (skip_ai_handoff). ****? \$0.3?0.7 total****, ~one L4 exec as projected ? but the stitch's real cost is ****concurrency**** (~16 min of control-plane heavy CPU), not dollars.

****Net:**** both owner-gated cells are now live-measured on rev ``00019`` ? validation wall ****183 s**** (corroborating the 2?3 min projection) and the reel compiles end-to-end. #512 one-pass validation and #513 skip-validation-on-real-dispatch are now both directly observed in prod. The Video-1 412 s?~170 s A/B remains a projection by design (different, heavier clip).

11. assistant (2026-07-08T05:05:49.890Z)

I have the full picture from S9. Now let me study the core architecture files. Let me read them in parallel.

12. user (2026-07-08T05:05:51.195Z)

```
1      """Processing / compile orchestration, extracted from ``backend.main`` (P23; re-cut of
#390).
2
3      The full hash -> validate -> segment -> report pipeline (`_execute_processing`), the
4      Cloud Run dispatch + ingest path (`_maybe_dispatch_remote` /
`_ingest_remote_segments`),
5      and the shared compile resolution + worker handler (`_resolve_compile` /
`_execute_compile`).
6      Extracted verbatim from ``backend.main`` to shrink the orchestrator; behavior is
unchanged.
7
8      SEAM DISCIPLINE (same convention as `backend.api.job_worker`): anything that must
remain
9      monkeypatchable-on-``backend.main`` or that lives in ``backend.main`` (the transitional
10     `db_instance`/`global_config` globals and the `ProcessRequest` model) is resolved
through
11     the `backend.main` module object at call time (`import backend.main as _main`),
never as a
12     bare local. `backend.main` re-exports every name below, so
`monkeypatch.setattr(backend.main,
13     `...)` and `import backend.main as m; m._execute_processing(...)` keep working
unchanged.
14     This module must NOT import `backend.main` at module level (main imports us for the
re-export).
15     Routed seams (rebind-patched on `backend.main` by tests): `global_config`,
`db_instance`,
16     `ProcessRequest`, `compute_video_id`, `emit_indexer_report`, `VideoCompiler`.
17     Object-attribute patches (`m.storage.fetch`, `m.segmenter_registry.get`, ...) hit
the shared
18     objects and need no routing.
19     """
20
21     import csv
22     import json
23     import logging
24     import math
25     import os
26     import re
27     import tempfile
28     import threading
29     from typing import TYPE_CHECKING, Any, Dict, List, Optional, Tuple
30     from urllib.parse import quote
31
32     from backend import storage
33     from backend.config import AppConfig, StitchingConfig
```

```

34     from backend.pipeline.segmenters.base import segmenter_registry
35     from backend.pipeline.stitcher.watermark_i18n import localize_watermark
36     from backend.pipeline.validators.base import ValidationResult, registry
37     from backend.results import Failure
38     from backend.utils.paths import output_base, safe_output_filename
39
40     if (
41         TYPE_CHECKING
42     ): # runtime import would be circular (main imports us for the re-export)
43         from backend.main import ProcessRequest # noqa: F401
44
45     logger = logging.getLogger(__name__)
46
47     # #466: serialize CPU/GPU-heavy LOCAL work ? in-process segmentation and the ffmpeg
compile stitch ?
48     # so that raising the drain's worker count (to overlap I/O-bound REMOTE Cloud Run
dispatch/poll) never
49     # accidentally parallelizes local work. A single self-hosted box serializes GPU work and
Gemini is
50     # rate-fragile; a remote dispatch is a bucket-poll and never holds this lane, so N
remote jobs overlap
51     # while local/compile jobs stay 1-at-a-time. It's a plain mutex (no nesting) ? no
deadlock; when the
52     # drain is serial (max_workers=1) it's an uncontended no-op = today's behaviour.
53     _LOCAL_COMPUTE_LANE = threading.BoundedSemaphore(1)
54
55
56     def _finalize_reingest(db, video_id: str, segments, play_wm: int, cand_wm: int) -> None:
57         """Commit a (re)segmentation with destroy-AFTER-confirm semantics.
58
59         On a non-empty run, drop the PRIOR rows (id <= the pre-run watermark) so only the
60         fresh segments remain. On an empty run, drop the NEW partial rows (id > watermark)
and
61         keep the prior good results intact ? a failed/empty re-run never destroys prior
data.
62         """
63         if segments:
64             db.prune_segments(video_id, play_upto=play_wm, cand_upto=cand_wm)
65         else:
66             db.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)
67
68
69     def _validation_failure_message(failures: List[Any]) -> str:
70         details = []
71         for failure in failures:
72             name = getattr(failure, "validator_name", None) or "validation"
73             message = (getattr(failure, "message", None) or "").strip()
74             details.append(f"{name}: {message}" if message else str(name))
75         if not details:
76             return "Video failed validation checks."
77         return "Video failed validation: " + "; ".join(details)
78
79
80     def _report_config_for_input(config: Any, raw_input: str) -> Any:
81         """Return report config that keeps remote-input reports out of disposable
scratch."""
82         storage_config = (
83             config.get("storage")
84             if isinstance(config, dict)
85             else getattr(config, "storage", None)
86         )
87         try:
88             if storage.input_is_local(raw_input, storage_config):
89                 return config
90         except ValueError:

```

```

91         return config
92     try:
93         cfg = (
94             config.model_dump(mode="json")
95             if hasattr(config, "model_dump")
96             else dict(config or {})
97         )
98     except Exception:
99         cfg = {}
100     report = dict(cfg.get("report") or {})
101     if report.get("output_dir"):
102         return config
103     report["output_dir"] = os.path.join(output_base(config), "reports")
104     cfg["report"] = report
105     return cfg
106
107
108     def _ingest_remote_segments(
109         db, video_id: str, outputs_uri: str, storage_config: dict, *, max_rows: int = 20000
110     ) -> int:
111         """Read a cloud job's ``intervals.csv`` (start,end,shot_count,ending_reason) from
112         the bucket and
113         persist each rally via the SAME ``db.add_play_segment`` the in-process segmenters
114         call ? so
115         ``/api/query`` + ``/api/compile`` read the rows with zero knowledge of where they
116         were produced.
117         Returns the count ingested. Mirrors ``cloud_run._read_json``'s fetch-to-a-real-dst
118         pattern (cloud
119         backends REQUIRE a dst). Hardened against a hostile/corrupt worker output: a
120         malformed row is
121         skipped, non-finite / overflowing values are rejected, the loop is bounded, and the
122         temp dir is
123         cleaned. May raise (missing/unreadable CSV, fetch error) ? the caller treats that as
124         a job failure
125         and prunes any partial rows (destroy-AFTER-confirm)."""
126         uri = f"{outputs_uri.rstrip('/')}/intervals.csv"
127         ref = storage.parse_uri(uri)
128         n = 0
129         with tempfile.TemporaryDirectory(prefix="cr-ingest-") as td:
130             local = storage.fetch(
131                 uri, os.path.join(td, ref.name or "intervals.csv"), config=storage_config
132             )
133             with open(local, newline="", encoding="utf-8") as f:
134                 for row in csv.DictReader(f):
135                     if n >= max_rows:
136                         logger.warning(
137                             "[API] cloud ingest hit the %d-row cap (%s) ? ignoring the
138                             rest.",
139                             max_rows,
140                             uri,
141                         )
142                         break
143                     try:
144                         start, end = float(row["start"]), float(row["end"])
145                     except (KeyError, TypeError, ValueError):
146                         continue
147                     if not (math.isfinite(start) and math.isfinite(end)):
148                         continue # reject inf/nan times rather than persist a junk rally
149                     meta: Dict[str, Any] = {
150                         "source": "cloud_run",
151                         "ending_reason": (row.get("ending_reason") or "").strip(),
152                     }
153                     sc = (row.get("shot_count") or "").strip()
154                     if sc:
155                         try:

```

```

148             meta["shot_count"] = int(
149                 float(sc)
150             ) # OverflowError on 'inf'/'1e400'
151         except (ValueError, OverflowError):
152             meta["shot_count"] = sc
153         db.add_play_segment(video_id, start, end, metadata=meta)
154         n += 1
155     return n
156
157
158     # Trajectory segmenters whose IN-PROCESS detect needs model weights present on THIS box.
A hosted
159     # cloud-serving control-plane bakes NO detector weights into its image (only the worker
Cloud Run Job
160     # GCSFuse-mounts them), so forcing compute="local" for one of these there can never
succeed ? the
161     # guardrail in _maybe_dispatch_remote fails such a request FAST instead of the cryptic
deep
162     # "WASB weights not found". (Names are the registered segmenter identities; see
backend/pipeline/segmenters.)
163     _WEIGHTS_BACKED_SEGMENTERS = frozenset({"wasb_hybrid", "tracknet_hybrid",
"fusion_hybrid"})
164
165
166     def _local_detector_weights_present(cfg, seg_name: str) -> bool:
167         """Whether the trajectory weights the IN-PROCESS ``seg_name`` detector needs are
present + readable
168         on THIS box. Mirrors the native/WSL runners' resolution (env override ?
``indexing.<det>.weights_path``)
169         and the ``api/capabilities`` probe, so the guardrail agrees with the healthcheck
that would
170         otherwise fail deep in the run. Fusion follows its configured substrate (WASB
default / TrackNet).
171         A non-weights segmenter returns True (nothing local to check)."""
172         name = (seg_name or "").strip().lower()
173         if name not in _WEIGHTS_BACKED_SEGMENTERS:
174             return True
175         idx = getattr(cfg, "indexing", None)
176         substrate = (
177             getattr(getattr(idx, "fusion", None), "substrate", "") or "wasb"
178         ).strip().lower()
179         if name == "tracknet_hybrid" or (name == "fusion_hybrid" and substrate ==
"tracknet"):
180             path = getattr(getattr(idx, "tracknet", None), "weights_path", "") or ""
181         else: # wasb_hybrid, or fusion on the (default) WASB substrate
182             path = (
183                 os.environ.get("WASB_WEIGHTS")
184                 or getattr(getattr(idx, "wasb", None), "weights_path", "")
185                 or ""
186             )
187         return bool(path) and os.path.exists(os.path.expanduser(path))
188
189
190     def _cloud_available(cfg) -> bool:
191         """Whether THIS server can satisfy a per-request ``compute="cloud"`` dispatch (item
3).
192
193         True when the control plane is wired for Cloud Run: either the configured default
target is
194         already ``cloud_run`` (the operator chose cloud serving), or the Cloud Run env
coordinates
195         (``CLOUD_RUN_JOB`` + ``GOOGLE_CLOUD_PROJECT``) and an output bucket
(``storage.output_root``)
196         are present so a forced per-request override can actually dispatch + collect a
result.

```

```

197
198     Surfaced in ``/api/capabilities`` so the SPA only offers the cloud toggle when it
would work."""
199     deployment = getattr(cfg, "deployment", None)
200     if (getattr(deployment, "compute_target", "") or "") == "cloud_run":
201         return True
202     has_job = bool(
203         os.environ.get("CLOUD_RUN_JOB") and os.environ.get("GOOGLE_CLOUD_PROJECT")
204     )
205     out_root = (getattr(getattr(cfg, "storage", None), "output_root", "") or "").strip()
206     return bool(has_job and out_root)
207
208
209 def _wait_progress_message(seg_name: str, elapsed_sec: float) -> str:
210     """The live job.progress line while a cloud GPU run is in flight (#482). Honest and
211     minute-granular: the control plane knows ELAPSED time, not frames ? frame-level %
comes
212     with the worker-side heartbeat (the #482 follow-up). Under a minute reads as
"starting"
213     (queue/cold-start), so the very first tick already shows movement."""
214     minutes = int(elapsed_sec) // 60
215     if minutes < 1:
216         return f"segmenting ({seg_name}) on cloud GPU ? starting up"
217     return f"segmenting ({seg_name}) on cloud GPU ? {minutes} min elapsed"
218
219
220 def submit_time_video_id(input_ref, cfg) -> Optional[str]:
221     """The video's md5 (the canonical ``video_id``) resolvable CHEAPLY at submit time ?
gs://
222     object metadata only, one stat, no bytes (#484). ``None`` ? unknown: local files
keep
223     deferred hashing on the drain (DEFERRED_HARDENING #16), and composite objects carry
no
224     ``md5Hash``. Never raises ? a metadata miss must not block a submit."""
225     if getattr(input_ref, "backend", "") != "gcs":
226         return None
227     sc = _storage_settings(cfg)
228     if not sc:
229         return None
230     try:
231         st = storage.get_storage("gcs", config=sc).stat(input_ref)
232         return getattr(st, "md5", None) or None
233     except Exception: # noqa: BLE001 ? best-effort identity, the drain still hashes
234         logger.info(
235             "[API] submit-time md5 unavailable for %s (metadata stat failed)",
236             getattr(input_ref, "uri", input_ref),
237             exc_info=True,
238         )
239     return None
240
241
242 def _cached_process_payload(
243     video_id: str, validation_results, segments, requested_compute
244 ) -> Dict[str, Any]:
245     """The /api/process result for a video whose detection already exists ? ONE shape
for
246     both producers (the drain's DB-cache short-circuit and the submit-time cache,
#484)."""
247     return {
248         "video_id": video_id,
249         "status": "success",
250         "message": (
251             f"Reused {len(segments)} cached play segments. "
252             "Pass force=true to reprocess this video."
253         ),

```

```

254         "results": validation_results or [],
255         "play_segments": segments,
256         "cached": True,
257         "compute": {"path": "cached", "requested": requested_compute},
258     }
259
260
261     def _read_bucket_json(uri: str, storage_config: dict) -> Optional[Dict[str, Any]]:
262         """Fetch + parse one small JSON object (cloud backends require a real dst ? the
263         ``_ingest_remote_segments`` pattern). Raises on fetch/parse errors; callers
264         decide."""
265         ref = storage.parse_uri(uri)
266         with tempfile.TemporaryDirectory(prefix="cr-report-") as td:
267             local = storage.fetch(
268                 uri, os.path.join(td, ref.name or "report.json"), config=storage_config
269             )
270             with open(local, encoding="utf-8") as f:
271                 loaded = json.load(f)
272             return loaded if isinstance(loaded, dict) else None
273
274     # One re-hydrate at a time per video_id (#492 review blocker 1): two concurrent cache
275     hits
276     # from an empty DB must not both ingest intervals.csv ? the loser waits, then finds the
277     # winner's rows in the DB and reuses them. Process-wide is enough: the CP is single-
278     instance
279     # (maxScale=1, spec D8) and SQLite serializes cross-process writers anyway.
280     _REHYDRATE_LOCKS: Dict[str, threading.Lock] = {}
281     _REHYDRATE_LOCKS_GUARD = threading.Lock()
282
283     def _rehydrate_lock(video_id: str) -> threading.Lock:
284         with _REHYDRATE_LOCKS_GUARD:
285             return _REHYDRATE_LOCKS.setdefault(video_id, threading.Lock())
286
287     def probe_process_cache(db, cfg, video_id: str) -> bool:
288         """READ-ONLY: does completed detection for ``video_id`` already exist (DB or
289         bucket)?
290
291         Safe to call before the atomic submit claim ? it writes nothing (#492 review blocker
292         1
293         ordering). Used to decide whether a submit can skip the local-scratch fetch guard
294         (blocker 2): a True probe means the request should never need CP scratch. Best-
295         effort:
296         errors report False (? the normal dispatch path with its guards)."""
297         try:
298             existing = db.get_video(video_id)
299             if (
300                 existing
301                 and existing.get("status") == "processed"
302                 and db.query_play_segments(video_id, {})
303             ):
304                 return True
305             sc = _storage_settings(cfg)
306             if not sc or (sc.get("backend") or "local") == "local":
307                 return False
308             output_root = (sc.get("output_root") or "").rstrip("/")
309             if not output_root:
310                 return False
311             report_uri = f"{output_root}/outputs/{video_id}/report.json"
312             if not storage.exists(report_uri, config=sc):
313                 return False
314             report = _read_bucket_json(report_uri, sc)
315             return bool(report and report.get("status") == "success")

```

```

313         except Exception: # noqa: BLE001 ? a failed probe degrades to the guarded dispatch
path
314             logger.warning(
315                 "[API] cache probe failed for %s ? treating as a miss", video_id,
exc_info=True
316             )
317             return False
318
319
320     def cached_process_result(
321         db, cfg, video_id: str, video_path: str, requested_compute=None
322     ) -> Optional[Dict[str, Any]]:
323         """A ready-to-serve /api/process result when this video's detection ALREADY exists ?
324         the DB first, then the bucket (``outputs/<md5>/report.json`` with ``status=success``
325         re-ingests via ``_ingest_remote_segments``): **the DB is a cache, the bucket is the
326         source of truth** (#484 instant re-runs; the #485 re-hydrate mechanism). ``None`` ?
no
327         completed work found (or the check failed) ? the caller dispatches normally. The
bucket
328         branch mirrors the real cloud ingest exactly (watermarks + ``_finalize_reingest``),
so a
329         partial earlier ingest is replaced, never double-counted.
330
331         Concurrency (#492 review blocker 1): the whole check-then-ingest runs under a
332         per-``video_id`` lock ? a concurrent caller blocks, then finds the winner's rows via
the
333         in-lock DB re-check and reuses them (one segment set, ever). Callers must hold a
CLAIMED
334         job before invoking (writes happen post-claim)."""
335         try:
336             with _rehydrate_lock(video_id):
337                 existing = db.get_video(video_id)
338                 if existing and existing.get("status") == "processed":
339                     segs = db.query_play_segments(video_id, {})
340                     if segs:
341                         return _cached_process_payload(
342                             video_id,
343                             existing.get("validation_results", []),
344                             segs,
345                             requested_compute,
346                         )
347                 sc = _storage_settings(cfg)
348                 if not sc or (sc.get("backend") or "local") == "local":
349                     return None
350                 output_root = (sc.get("output_root") or "").rstrip("/")
351                 if not output_root:
352                     return None
353                 outputs_uri = f"{output_root}/outputs/{video_id}"
354                 report_uri = f"{outputs_uri}/report.json"
355                 if not storage.exists(report_uri, config=sc):
356                     return None
357                 report = _read_bucket_json(report_uri, sc)
358                 if not report or report.get("status") != "success":
359                     return None
360                 if not video_path:
361                     # Read-path re-hydrate (#485): a wiped DB has no source record. The
report
362                     # carries the worker's --video-uri (P3); older reports fall back to the
363                     # staged-canonical videos/<md5>/ prefix; else "" ? the segments still
364                     # recover, and the next same-md5 submit re-binds the truthful path via
365                     # add_video's upsert.
366                     video_path = str(report.get("video_uri") or "") or _first_object_under(
367                         f"{output_root}/videos/{video_id}/", sc
368                     )
369                 play_wm, cand_wm = db.segment_watermarks(video_id)

```



```

370         db.add_video(video_id, video_path, "processed", [])
371         n = _ingest_remote_segments(db, video_id, outputs_uri, sc)
372         if n <= 0:
373             return None
374         segments = [
375             s
376             for s in db.query_play_segments(video_id, {})
377             if int(s.get("id", 0)) > play_wm
378         ]
379         _finalize_reingest(db, video_id, segments, play_wm, cand_wm)
380         if not segments:
381             return None
382         logger.info(
383             "[API] re-hydrated %d segments for video_id=%s from %s (submit-time
cache hit)",
384             n,
385             video_id,
386             outputs_uri,
387         )
388         return _cached_process_payload(video_id, [], segments, requested_compute)
389     except Exception: # noqa: BLE001 ? cache misses must degrade to a normal dispatch
390         logger.warning(
391             "[API] submit-time cache check failed for %s ? dispatching normally",
392             video_id,
393             exc_info=True,
394         )
395     return None
396
397
398     def _first_object_under(prefix_uri: str, storage_config: dict) -> str:
399         """The first object URI under a bucket prefix, or "" ? best-effort source
recovery."""
400         try:
401             return next(iter(storage.list_uris(prefix_uri, config=storage_config)), "")
402         except Exception: # noqa: BLE001 ? recovery nicety, never a failure
403             return ""
404
405
406     _MD5_HEX_RE = re.compile(r"[0-9a-f]{32}")
407
408
409     def read_path_rehydrate(db, cfg, video_id: str) -> bool:
410         """#485, the read half of "DB = cache, bucket = truth": before a rally READ reports
411         nothing for an md5-shaped id, try re-hydrating the video from ``outputs/<md5>/` ` ?
so
412         ``?video=<md5>`` lands on grounded state after any restart/redeploy wipe. True ? the
DB
413         has the video's rows now (re-query). Non-md5 ids return False without touching the
414         bucket; every failure degrades to False (the read just reports what the DB has)."""
415         vid = (video_id or "").strip().lower()
416         if not _MD5_HEX_RE.fullmatch(vid):
417             return False
418         if not probe_process_cache(db, cfg, vid):
419             return False
420         return cached_process_result(db, cfg, vid, "") is not None
421
422
423     def _maybe_dispatch_remote(
424         cfg,
425         db,
426         req: "ProcessRequest",
427         video_id: str,
428         seg_name: str,
429         val_results,
430         play_wm: int,

```

```

431         cand_wm: int,
432         notify,
433     ) -> Optional[Tuple[Dict[str, Any], bool]]:
434         """Cloud-serving (COMPUTE_DECOUPLED_SERVING M6 ? the API path): when
435         ``deployment.compute_target
436         == "cloud_run"`, run the heavy DETECT on a Cloud Run GPU Job and INGEST its rally
437         intervals into
438         THIS DB, so the rest of the flow (query ? compile/local-stitch) is unchanged.
439
440         Returns ``(response_dict, ran)`` when the cloud path is taken ? ``ran`` is whether
441         the cloud job
442         actually EXECUTED (``False`` for a pre-dispatch config rejection, so the
443         observability report
444         doesn't claim a fake segmentation run) ? or ``None`` to fall through to the in-
445         process segmenter
446         (**DEFAULT-OFF*: ``get_compute_backend`` returns ``None`` for every
447         non-``cloud_run`` target, so
448         the in-process path stays byte-identical)."""
449         from backend.compute import ComputeJobSpec, get_compute_backend
450
451         def _failed(msg: str, ran: bool) -> Tuple[Dict[str, Any], bool]:
452             # Shape-match the in-process segmenter-fail branch + drop any partial NEW rows
453             (id > play_wm);
454             # prior good rows (id <= play_wm) are preserved (destroy-AFTER-confirm).
455             db.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)
456             # Provenance even on failure, and HONEST about what ran: `path` is "cloud_run"
457             only when the
458             # job actually dispatched + executed remotely (ran=True, e.g. a non-zero worker
459             rc / ingest
460             # error); a pre-dispatch reject (ran=False) ran NOWHERE ? "not_run". `target`
461             records the
462             # intended backend, `requested` the picker choice, `dispatched` whether the
463             remote job ran.
464             return (
465                 {
466                     "video_id": video_id,
467                     "status": "failed",
468                     "message": msg,
469                     "results": [r.model_dump() for r in val_results],
470                     "play_segments": [],
471                     "compute": {
472                         "path": "cloud_run" if ran else "not_run",
473                         "requested": req.compute,
474                         "target": "cloud_run",
475                         "dispatched": ran,
476                     },
477                 },
478                 ran,
479             )
480
481         # SPA compute-picker (item 3): a per-request choice overrides the server's
482         configured default.
483         # "local" forces the in-process segmenter; "cloud" forces a Cloud Run dispatch
484         (guarded by
485         # _cloud_available so we fail clearly rather than silently running local); None /
486         unknown ? default.
487         choice = (req.compute or "").strip().lower()
488         if choice and choice not in ("local", "cloud"):
489             logger.warning(
490                 "[API] ignoring unknown compute=%r (expected 'local'/'cloud'); using server
491                 default.",
492                 req.compute,
493             )
494             choice = ""
495         if choice == "local":

```

```

481         # A hosted cloud-serving control-plane bakes NO detector weights into its image
(only the
482         # worker Cloud Run Job GCSFuse-mounts them), so forcing a weights-backed
trajectory detector
483         # to run IN-PROCESS here can never succeed ? it dies deep in the segmenter with
a cryptic
484         # "WASB weights not found" (observed live: job f29d41d0?, 2026-07-07). Fail FAST
with an
485         # actionable message instead of that silent trap. A truly-local box (not cloud-
capable), or a
486         # cloud-wired box that ACTUALLY has the weights on disk, is unaffected ? it runs
in-process
487         # below exactly as before (the weights-present check is the escape hatch).
488         if (
489             (seg_name or "").strip().lower() in _WEIGHTS_BACKED_SEGMENTERS
490             and _cloud_available(cfg)
491             and not _local_detector_weights_present(cfg, seg_name)
492         ):
493             return _failed(
494                 f"this hosted server can't run detection locally ? the '{seg_name}'
model weights "
495                 "live only on the cloud worker, not on the control-plane. Choose 'Run in
cloud'.",
496                 False,
497             )
498         return None # explicit "run on my machine" ? in-process regardless of the
server's target
499         if choice == "cloud":
500             if not _cloud_available(cfg):
501                 return _failed(
502                     "cloud-serving was requested but this server isn't configured for it "
503                     "(no Cloud Run target). Pick 'run on my machine', or configure cloud
serving.",
504                     False,
505                 )
506             compute = get_compute_backend(cfg, target_override="cloud_run")
507         else:
508             compute = get_compute_backend(
509                 cfg
510             ) # server default (default-OFF unless cloud_run)
511         if compute is None:
512             return None # default-OFF ? run the in-process segmenter below (unchanged)
513
514         storage_cfg = getattr(cfg, "storage", None)
515
516         # The Cloud Run job fetches the input FROM THE BUCKET ? it can't read a path local
to this box.
517         try:
518             input_ref = storage.resolve_input_ref(req.video_path, storage_cfg)
519         except ValueError as e:
520             return _failed(
521                 f"cloud-serving: unresolvable input '{req.video_path}': {e}", False
522             )
523         if input_ref.backend == "local":
524             # A browser upload lands on THIS box's local disk; the Cloud Run job (a
different machine)
525             # can't read it. Stage it UP to the input bucket first, then dispatch from there
(#461). Only
526             # possible when a CLOUD input_root is configured; without one we genuinely can't
stage ? fail.
527             # Classify the ROOT ITSELF (parse_uri), not storage.input_backend: the staged
URI below is
528             # built by joining under input_root, and storage.put routes by
parse_uri(staged_uri).backend ?
529             # so an explicit cloud root (gs://?, s3://?) stages fine with input_backend left

```

```

at its
530         # default, while a bare/local-looking root can't stage no matter what
input_backend says
531         # (resolve_input_uri only consults input_backend when input_root is EMPTY, and
an empty root
532         # gives us nothing to build the staged URI under).
533         input_root = (getattr(storage_cfg, "input_root", "") or "").strip()
534         try:
535             root_backend = (
536                 storage.parse_uri(input_root).backend if input_root else "local"
537             )
538         except ValueError: # unsupported scheme in input_root ? nowhere real to stage
to
539             root_backend = "local"
540         if root_backend == "local":
541             return _failed(
542                 "cloud-serving needs a cloud input (e.g. gs://?), or a cloud
storage.input_root to "
543                 "stage local uploads into; a path local to this machine can't be read by
the Cloud "
544                 "Run job.",
545                 False,
546             )
547         local_src = input_ref.key # resolved local path of the upload
548         staged_uri =
f"{input_root.rstrip('/')}/{video_id}/{os.path.basename(local_src)}"
549         try:
550             notify("staging upload to bucket")
551             # Thread the active storage settings so backend-specific config survives (S3
endpoint_url/
552             # region/addressing for R2·D0·MinIO; GCS project/client) ? not a fresh
default client.
553             storage.put(
554                 local_src,
555                 staged_uri,
556                 config=(
557                     storage_cfg.model_dump() if storage_cfg is not None else None
558                 ),
559             )
560         except Exception as e: # noqa: BLE001 ? a staging failure is a clean job
failure, not a 500
561             logger.error(
562                 "[API] cloud-serving: failed to stage %s -> %s: %s",
563                 local_src,
564                 staged_uri,
565                 e,
566                 exc_info=True,
567             )
568             return _failed(
569                 f"cloud-serving: failed to stage the uploaded file to {input_root}:
{e}",
570                 False,
571             )
572         # Repoint BOTH the dispatch and the DB video record at the staged bucket object:
the L4 job
573         # DETECTs from it, and a later /api/compile re-fetches the same source (the
local upload is
574         # gone once this instance recycles). add_video is an UPSERT ? it only updates
filepath.
575         req.video_path = staged_uri
576         input_ref = storage.resolve_input_ref(staged_uri, storage_cfg)
577         db.add_video(
578             video_id, staged_uri, "processed", [r.model_dump() for r in val_results]
579         )
580         out_root = (getattr(storage_cfg, "output_root", "") or "").strip()

```

```

581         if not out_root:
582             return _failed(
583                 "cloud-serving needs storage.output_root set to a bucket (e.g. gs://your-
bucket) "
584                 "? that's where the Cloud Run job writes the result.",
585                 False,
586             )
587
588         notify("dispatching (cloud_run)")
589         logger.info(
590             "[API] cloud-serving: dispatching %s -> Cloud Run Job (out=%s)",
591             req.video_path,
592             out_root,
593         )
594         # Forward the engine's AI-handoff stance so the cloud job runs the SAME detection:
with
595         # skip_ai_handoff=True the WASB windows become rallies directly (no Gemini key
needed); without
596         # forwarding the cloud job would default to the handoff and yield 0 rallies when it
has no key.
597         # (player_pool / max_frames are still dropped on the cloud path ? a separate
documented follow-up.)
598         overrides = []
599         idx_cfg = getattr(cfg, "indexing", None)
600         sk = getattr(idx_cfg, "skip_ai_handoff", None)
601         if sk is not None:
602             overrides.append(f"indexing.skip_ai_handoff={'true' if sk else 'false'}")
603         # Forward the resolved WASB serving knobs so the PRESET governs the worker
(#506/#507). The
604         # worker Job has no DEPLOYMENT_PROFILE, so it does NOT apply the cloud-serving
preset itself ?
605         # without forwarding, its decode_backend + GPU-feed tuning silently fall back to the
model
606         # defaults (cpu, batch 8) regardless of what the control-plane resolved.
decode_backend is
607         # cuda-only (validator-enforced) and requires a worker image that understands
--decode-backend
608         # (? #507); batch_size/prefetch_batches are forwarded only when set (None = the
parity default).
609         wasb_cfg = getattr(idx_cfg, "wasb", None)
610         if wasb_cfg is not None:
611             decode_backend = getattr(wasb_cfg, "decode_backend", None)
612             if decode_backend:
613                 overrides.append(f"indexing.wasb.decode_backend={decode_backend}")
614             batch_size = getattr(wasb_cfg, "batch_size", None)
615             if batch_size:
616                 overrides.append(f"indexing.wasb.batch_size={batch_size}")
617             prefetch = getattr(wasb_cfg, "prefetch_batches", None)
618             if prefetch:
619                 overrides.append(f"indexing.wasb.prefetch_batches={prefetch}")
620         # Tell the worker to SKIP its redundant re-validation: THIS control-plane is the
authoritative
621         # gate ? it already ran the validators with its resolved config just above (only a
PASS reaches
622         # here; a failure returned "failed" before dispatch). A worker re-run would decode
the video a
623         # second time and, because its config isn't fully forwarded, can resolve DIFFERENT
thresholds and
624         # REJECT a clip this gate already passed (observed 2026-07-07: passed at
min_blur_threshold=8.0,
625         # then worker-rejected at its default 20.0 after a successful GPU dispatch).
626         spec = ComputeJobSpec(
627             video_uri=req.video_path,
628             out_uri=out_root,
629             segmenter=seg_name,

```

```

630         config_overrides=tuple(overrides),
631         skip_validation=True,
632     )
633     try:
634         # Dispatch the Cloud Run Job + bucket-poll to completion. on_wait is the live
635         # heartbeat (#482): job.progress moves on every poll tick instead of freezing at
636         # "segmenting" for the whole GPU run (the 2026-07-05 silent 22-min spinner) ?
637         # also gives the SPA a stall signal to time out on, instead of a blind wall
638         # clock.
639         result = compute.run_infer(
640             spec,
641             on_wait=lambda elapsed: notify(_wait_progress_message(seg_name, elapsed)),
642         )
643     except Exception as e: # noqa: BLE001 ? any dispatch/poll error is a clean job
644         failure, not a 500
645         logger.error("[API] cloud dispatch error: %s", e, exc_info=True)
646         return _failed(f"cloud dispatch error: {e}", True)
647     if result.returncode != 0:
648         return _failed(f"cloud job failed (returncode {result.returncode}).", True)
649     if result.video_id and result.video_id != video_id:
650         # The worker re-hashes the bucket bytes; a divergence means the input changed
651         # between this
652         # box's hash and the worker's ? ingest under the LOCAL id but flag it so it's
653         # diagnosable.
654         logger.warning(
655             "[API] cloud video_id %s != local %s ? input bytes differ between the API
656             hash "
657             "and the worker hash; ingesting under the local id.",
658             result.video_id,
659             video_id,
660         )
661     notify("ingesting (cloud_run)")
662     storage_dict = (
663         storage_cfg.model_dump() # type: ignore[union-attr] # hasattr-guarded; mypy
664         can't narrow
665         if hasattr(storage_cfg, "model_dump")
666         else {}
667     )
668     try:
669         _ingest_remote_segments(db, video_id, result.outputs_uri, storage_dict)
670     except Exception as e: # noqa: BLE001 ? a missing/partial/corrupt result must not
671         500 or leak rows
672         logger.error("[API] cloud ingest failed: %s", e, exc_info=True)
673         return _failed(
674             f"cloud-serving: could not read the result from {result.outputs_uri}: {e}",
675             True,
676         )
677     # Only the NEW cloud rows (id > the pre-run watermark); _finalize_reingest then
678     # drops the old ones.
679     segments = [
680         s for s in db.query_play_segments(video_id, {}) if int(s.get("id", 0)) > play_wm
681     ]
682     _finalize_reingest(db, video_id, segments, play_wm, cand_wm)
683     # Compute-path PROVENANCE: the GCP-verifiable proof the cloud path ran. `execution`
684     # is the real
685     # Cloud Run execution name (cross-check: gcloud run jobs executions describe
686     # <execution>).
687     provenance = {
688         "path": "cloud_run",
689         "requested": req.compute,
690         "job": os.environ.get("CLOUD_RUN_JOB", "rally-worker"),
691         "region": os.environ.get("CLOUD_RUN_REGION", "asia-south1"),

```

```

684         "execution": result.execution,
685         "outputs_uri": result.outputs_uri,
686     }
687     logger.info(
688         "[COMPUTE] cloud_run dispatched: execution=%s job=%s region=%s out=%s
video_id=%s",
689         result.execution,
690         provenance["job"],
691         provenance["region"],
692         result.outputs_uri,
693         video_id,
694     )
695     return (
696         {
697             "video_id": video_id,
698             "status": "success",
699             "message": f"Successfully indexed {len(segments)} play segments (cloud_run /
'{seg_name}').",
700             "results": [r.model_dump() for r in val_results],
701             "play_segments": segments,
702             "compute": provenance,
703         },
704         True,
705     )
706
707
708 def _execute_processing(
709     config=None, db=None, req=None, progress=None
710 ) -> Dict[str, Any]:
711     """The full hash ? validate ? segment ? report pipeline for one video.
712
713     Accepts explicit ``config``/``db`` (the new app-factory pattern) OR reads the
714     module-level ``global_config``/``db_instance`` (backward-compatible with existing
715     tests that pass only ``req``). Route handlers always pass explicit params; older
716     test code that calls ``_execute_processing(req)`` still works.
717
718     ``progress`` is an optional callable(stage: str) used to surface coarse job
719     progress.
720
721     Raises on unexpected internal errors ? the caller records those as a job failure
722     """
723     import backend.main as _main # call-time seam: main's globals/model stay patchable
724
725     # Backward-compatible shim: if called as _execute_processing(req) (old pattern),
726     # shift positional args ? req arrives in position 0 (config), progress in position 1
727     (db).
728
729     if config is not None and not isinstance(config, AppConfig):
730         # Old signature: _execute_processing(req, progress=None)
731         if req is None and db is not None and isinstance(db, _main.ProcessRequest):
732             req = db
733             db = None
734         elif req is None and isinstance(config, _main.ProcessRequest):
735             req = config
736             config = None
737
738     # Resolve config/db from module globals when not passed explicitly
739     if config is None:
740         config = _main.global_config
741     if db is None:
742         db = _main.db_instance
743
744     notify = progress or (lambda stage: None)
745
746     notify("hashing")
747     # M2: resolve the input to a LOCAL path. A bare/local path is returned unchanged
748     (identity ?
749     # byte-identical to today); a bucket/Drive URI is fetched to scratch first. The

```

```

original
744     # req.video_path (the source URI) is kept for the DB record + report; only the file-
reading
745     # consumers (hash / validators / segmenter) use the resolved local path.
746     local_video = ""
747     try:
748         local_video = storage.acquire_local_input(
749             req.video_path, getattr(config, "storage", None)
750         )
751         video_id = _main.compute_video_id(local_video)
752     except Exception:
753         storage.cleanup_acquired_local_input(
754             req.video_path, local_video, getattr(config, "storage", None)
755         )
756         raise
757     try:
758         existing_video = db.get_video(video_id)
759     except Exception:
760         storage.cleanup_acquired_local_input(
761             req.video_path, local_video, getattr(config, "storage", None)
762         )
763         raise
764     was_reingest = existing_video is not None
765
766     # typed AppConfig: attribute access for the default segmenter (with safe fallback)
767     default_seg = getattr(
768         getattr(config, "indexing", None), "default_segmenter", "motion"
769     )
770     seg_name = req.segmenter or default_seg
771     segmenter_actual = None
772     segmentation_ran = False
773     val_results: List[Any] = []
774     val_context: Dict[str, Any] = {}
775     # Compute-path PROVENANCE (the validation method): which backend actually ran the
DETECT ?
776     # "in_process" once the local segmenter is reached, set on the response as
"cloud_run" by
777     # _maybe_dispatch_remote when it dispatches. None ? nothing ran (validation/config
fail).
778     compute_actual: Optional[str] = None
779     response: Optional[Dict[str, Any]] = None
780     try:
781         if (
782             existing_video
783             and existing_video.get("status") == "processed"
784             and not req.force
785         ):
786             cached_segments = db.query_play_segments(video_id, {})
787             if cached_segments:
788                 logger.info(
789                     "[API] Reusing cached segmentation for video_id=%s; force=true
bypasses cache.",
790                     video_id,
791                 )
792             response = _cached_process_payload(
793                 video_id,
794                 existing_video.get("validation_results", []),
795                 cached_segments,
796                 getattr(req, "compute", None),
797             )
798             return response
799
800     # 1. Run pluggable validators (with-context variant surfaces ffprobe_meta for
the report) ?
801     #     UNLESS this exact content already has a cached verdict under the current

```



```

validator config.
802         notify("validating")
803         # VALIDATION-RESULT CACHE: video_id is the MD5 of the file bytes, so an
identical-content
804         # reupload keys the same row. When that content was ALREADY validated under the
SAME validator
805         # configuration, reuse the stored verdict instead of re-decoding the clip
through the whole
806         # validator suite ? minutes on a heavy 4K clip on the CPU-only control-plane,
which the
807         # 2026-07-07 CUJ re-paid on every retry after a segmentation failure (validation
passed but
808         # detection failed ? no cached segments, so the whole-pipeline cache above
misses and we land
809         # here). The fingerprint folds in the validation.* config + sport + validator
set + a schema
810         # version, so ANY config change misses and re-validates (never serve a stale
PASS/FAIL).
811         # ``force=true`` bypasses the reuse and re-writes a fresh verdict below (mirrors
how it bypasses
812         # the whole-pipeline cache). Composes with #513 (the worker skips its own re-
validation on a
813         # cloud dispatch), so a post-failure reupload reaches segmentation in ~0
validation time.
814         val_fingerprint = registry.config_fingerprint(config)
815         # The whole read ? fingerprint-match ? replay is guarded: ANY cache fault
(unreadable row,
816         # corrupt or schema-incompatible stored results) degrades to a real validation
rather than
817         # failing the job. The cache is a pure speedup ? never a new failure mode.
818         reused_validation = False
819         if not req.force:
820             try:
821                 cached = db.get_validation_cache(video_id)
822                 if cached is not None and cached.get("fingerprint") == val_fingerprint:
823                     # Replay the stored ValidationResult list; val_context stays {} (no
fresh
824                     # ffprobe_meta ? the report's media block degrades gracefully to
size-only, the
825                     # deliberate cost of skipping the decode). Segmentation reads
neither of these.
826                     val_results = [
827                         ValidationResult(**r) for r in cached.get("results", [])
828                     ]
829                     reused_validation = True
830                     logger.info(
831                         "[API] Reusing cached validation verdict (passed=%s) for
video_id=%s; "
832                         "validator config unchanged ? skipping the validator suite. "
833                         "force=true re-validates.",
834                         cached.get("passed"),
835                         video_id,
836                     )
837             except Exception: # noqa: BLE001 ? any cache read/replay fault degrades to
a real validation
838                 reused_validation = False
839                 logger.warning(
840                     "[API] validation-cache reuse failed for video_id=%s ? re-
validating.",
841                     video_id,
842                     exc_info=True,
843                 )
844         if not reused_validation:
845             logger.info(f"Running validations for {req.video_path}...")
846             val_results, val_context = registry.run_all_with_context(local_video,

```

```

config)
847     failures = [r for r in val_results if not r.passed]
848     if not reused_validation:
849         # Persist the verdict keyed by (content, config) so the NEXT identical
reupload skips the
850         # decode. Best-effort (the writer swallows + logs its own faults); a miss
just re-validates.
851         db.set_validation_cache(
852             video_id,
853             val_fingerprint,
854             not failures,
855             [r.model_dump() for r in val_results],
856         )
857
858     if failures:
859         # add_video is an UPSERT (no cascade) ? a failed re-validation only updates
860         # status; it no longer destroys the video's prior good segments.
861         db.add_video(
862             video_id,
863             req.video_path,
864             "failed",
865             [r.model_dump() for r in val_results],
866         )
867         response = {
868             "video_id": video_id,
869             "status": "failed",
870             "message": _validation_failure_message(failures),
871             "results": [r.model_dump() for r in val_results],
872         }
873         return response
874
875     # 2. Run segmenter & indexer
876     logger.info("[API] Video passed checks. Segmenting and indexing...")
877     db.add_video(
878         video_id, req.video_path, "processed", [r.model_dump() for r in val_results]
879     )
880
881     segmenter_cls = segmenter_registry.get(seg_name)
882     if not segmenter_cls:
883         # Submit-time validation already 400s unknown names; this is the defensive
884         # in-worker path (e.g. config default pointing at an unregistered
segmenter).
885         available = list(segmenter_registry.list_available().keys())
886         response = {
887             "video_id": video_id,
888             "status": "failed",
889             "message": f"Segmenter '{seg_name}' not found. Available segmenters:
{available}",
890             "results": [r.model_dump() for r in val_results],
891             "play_segments": [],
892         }
893         return response
894
895     logger.info(f"[API] Using segmenter: {seg_name}")
896     notify(f"segmenting ({seg_name})")
897     segmenter_actual = seg_name
898     # Destroy-AFTER-confirm: remember the pre-run rows, then keep new + drop old
only
899     # if the run yields segments; otherwise drop the new partial rows and keep the
old.
900     play_wm, cand_wm = db.segment_watermarks(video_id)
901     # COMPUTE_DECOUPLED_SERVING M6 ? API: with deployment.compute_target=cloud_run,
run the heavy
902     # DETECT on a Cloud Run GPU Job + ingest its rallies HERE so query/compile are
unchanged.

```

```

903         # DEFAULT-OFF: returns None for every non-cloud target ? the in-process
segmenter below runs
904         # byte-identically. (player_pool/max_frames aren't threaded to the worker
`infer` CLI yet, so
905         # they're dropped on the cloud path in this slice ? a documented follow-up.)
906         remote = _maybe_dispatch_remote(
907             config, db, req, video_id, seg_name, val_results, play_wm, cand_wm, notify
908         )
909         if remote is not None:
910             response, segmentation_ran = (
911                 remote # `ran` = did the cloud job actually execute (report)
912             )
913             return response # the cloud branch already stamped response["compute"]
(path=cloud_run)
914         # In-process from here on ? record the actual path for the provenance stamp in
`finally`.
915         compute_actual = "in_process"
916         segmenter = segmenter_cls(db, config)
917         try:
918             # Serialize the CPU/GPU-heavy in-process detection (see
_LOCAL_COMPUTE_LANE). A remote
919             # dispatch never reaches here (it returned above), so remote polls still
overlap freely.
920             with _LOCAL_COMPUTE_LANE:
921                 result = segmenter.process_video(
922                     video_id, local_video, req.player_pool, req.max_frames
923                 )
924             except Exception:
925                 # A raising segmenter must not leave half-written rows: restore the pre-run
926                 # state (same destroy-after-confirm contract as the Failure branch), then
let
927                 # the job worker record the failure.
928                 db.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)
929                 raise
930             segmentation_ran = True
931
932             if isinstance(result, Failure):
933                 logger.error(f"[API] Segmenter failed: {result.message}")
934                 db.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)
935                 response = {
936                     "video_id": video_id,
937                     "status": "failed",
938                     "message": f"Segmenter '{seg_name}' failed: {result.message}",
939                     "results": [r.model_dump() for r in val_results],
940                     "play_segments": [],
941                 }
942                 return response
943
944             segments = result.value if hasattr(result, "value") else result
945             notify("finalizing")
946             _finalize_reingest(db, video_id, segments, play_wm, cand_wm)
947
948             response = {
949                 "video_id": video_id,
950                 "status": "success",
951                 "message": f"Successfully indexed {len(segments)} play segments using
'{seg_name}' segmenter.", # type: ignore[arg-type] # hasattr-unwrap above; Success arm has
.value
952                 "results": [r.model_dump() for r in val_results],
953                 "play_segments": segments,
954             }
955             return response
956         finally:
957             # Compute-path PROVENANCE stamp (the validation method): the cloud branch
already set

```

```

958         # response["compute"] (path=cloud_run + the Cloud Run execution); for every
other path stamp
959         # the actual path here so the job result PROVES which backend ran (in_process /
not_run) and
960         # what the picker REQUESTED ? a silent fallback (requested cloud, ran local) is
then visible.
961         if isinstance(response, dict):
962             response.setdefault(
963                 "compute",
964                 {
965                     "path": compute_actual or "not_run",
966                     "requested": getattr(req, "compute", None),
967                 },
968             )
969         prov = response.get("compute", {}) if isinstance(response, dict) else {}
970         logger.info(
971             "[COMPUTE] video_id=%s requested=%s actual=%s dispatched=%s status=%s",
972             video_id,
973             getattr(req, "compute", None),
974             prov.get("path"),
975             prov.get("dispatched"),
976             (response or {}).get("status") if isinstance(response, dict) else None,
977         )
978         # Always emit the per-video observability report (next to the video, or to
979         # report.output_dir). Never raises. Surface the paths in the response. (The
compute-path
980         # proof lives on the job result + the [COMPUTE] log, not the report ?
obsreport's record_run
981         # doesn't surface arbitrary run signals.)
982         paths = _main.emit_indexer_report(
983             db=db,
984             video_id=video_id,
985             # Resolved LOCAL path, not the raw URI: obsreport writes the report next to
the
986             # video via os.path.abspath(video_path), and a remote URI (gs://, gcs://,
local://)
987             # collapses to a literal scheme dir at the repo root (see #397). The source
URI is
988             # still preserved in DB provenance via db.add_video.
989             video_path=local_video,
990             config=_report_config_for_input(config, req.video_path),
991             val_results=val_results,
992             val_context=val_context,
993             segmenter_requested=req.segmenter,
994             segmenter_actual=segmenter_actual,
995             was_reingest=was_reingest,
996             segmentation_ran=segmentation_ran,
997         )
998         if paths and isinstance(response, dict):
999             response["reports"] = paths
1000         storage.cleanup_acquired_local_input(
1001             req.video_path, local_video, getattr(config, "storage", None)
1002         )
1003
1004
1005     def _remote_input_size_bytes(input_ref, storage_config) -> Optional[int]:
1006         """Best-effort size (bytes) of a remote storage source, or ``None`` if unknowable.
1007
1008         Used only to pre-reserve scratch for a fetch-to-local (P2). Never raises: a stat
that needs a
1009         missing cloud SDK, the network, or a credential returns ``None`` so the free-space
guard degrades
1010         to a no-op rather than blocking a job it cannot measure."""
1011         try:
1012             sc = (

```

```

1013         storage_config.model_dump()
1014         if hasattr(storage_config, "model_dump")
1015         else (storage_config or {})
1016     )
1017     backend = storage.get_storage(input_ref.backend, config=sc)
1018     return int(backend.stat(input_ref).size)
1019 except Exception as e: # noqa: BLE001 - measuring is best-effort; inability must
not block
1020     logger.debug(
1021         "free-space: could not stat remote source %s: %s", input_ref.uri, e
1022     )
1023     return None
1024
1025
1026 class _CompileError(Exception):
1027     """A submit-time-validatable compile problem ? carries the HTTP status + detail so
the API path
1028     maps it to an HTTPException (fast 4xx) and the worker path maps it to a failed-job
result."""
1029
1030     def __init__(self, status_code: int, detail: str):
1031         super().__init__(detail)
1032         self.status_code = status_code
1033         self.detail = detail
1034
1035
1036     def _resolve_compile(
1037         db, cfg, video_id: str, segment_ids: List[int], output_filename: str
1038     ):
1039         """Shared compile resolution ? run at SUBMIT (fast-fail 4xx) AND in the WORKER (re-
resolve, since a
1040         row could change between submit and run). Verifies the video exists, the requested
segments match
1041         and survive the ``stitching.min_shot_count`` floor (segments WITHOUT a known
shot_count are exempt
1042         so manual/legacy picks can't silently vanish), and the output name is traversal-
safe. Raises
1043         ``_CompileError(404/400)``. Returns ``(video_row, kept_segments,
safe_out_name)``."""
1044         video = db.get_video(video_id)
1045         if not video:
1046             raise _CompileError(404, "Indexed video record not found.")
1047         # Reject path traversal / absolute names (os.path.join would otherwise write
anywhere on disk).
1048         try:
1049             out_name = safe_output_filename(output_filename)
1050         except ValueError as e:
1051             raise _CompileError(400, str(e)) from e
1052
1053         all_segments = db.query_play_segments(video_id, {})
1054         matched = [s for s in all_segments if s["id"] in segment_ids]
1055         if not matched:
1056             raise _CompileError(400, "No matching play segments found to stitch.")
1057
1058         stitching = getattr(cfg, "stitching", None) or StitchingConfig()
1059         kept = []
1060         for s in matched:
1061             meta = s.get("metadata") or {}
1062             sc = meta.get("shot_count") if isinstance(meta, dict) else None
1063             try:
1064                 if sc is not None and int(sc) < stitching.min_shot_count:
1065                     continue
1066             except (TypeError, ValueError):
1067                 pass # unparseable count -> treat as unknown, keep
1068             kept.append(s)

```

```

1069         if len(kept) < len(matched):
1070             logger.info(
1071                 "[compile] stitching.min_shot_count=%s: dropped %d/%d segments below the
floor.",
1072                 stitching.min_shot_count,
1073                 len(matched) - len(kept),
1074                 len(matched),
1075             )
1076         if not kept:
1077             raise _CompileError(
1078                 400,
1079                 f"All {len(matched)} matching segments fall below "
1080                 f"stitching.min_shot_count={stitching.min_shot_count}.",
1081             )
1082         return video, kept, out_name
1083
1084
1085     def reel_object_uri(cfg, video_id: str, name: str) -> str:
1086         """The bucket URI a compiled reel is mirrored to for a direct signed-URL download
(#463), or "" when
1087         the output isn't a cloud bucket (truly-local ? serve the local file). Shared by the
compile
1088         mirror-upload and the /api/highlights redirect:
`<output_root>/highlights/<video_id>/<name>`. """
1089         root = (getattr(getattr(cfg, "storage", None), "output_root", "") or "").strip()
1090         if not root or "://" not in root:
1091             return ""
1092         return f"{root.rstrip('/')}/highlights/{video_id}/{name}"
1093
1094
1095     def _storage_settings(cfg) -> Optional[Dict[str, Any]]:
1096         """The active storage settings dict (per-backend config for the storage layer ? S3
endpoint/region,
1097         GCS project/client/signed_url_ttl), or None. Threaded into storage calls so a non-
default backend
1098         isn't silently constructed as a fresh default. """
1099         sc = getattr(cfg, "storage", None)
1100         return sc.model_dump() if sc is not None else None
1101
1102
1103     def signed_reel_url(cfg, video_id: str, name: str) -> Optional[str]:
1104         """A short-lived signed URL for a reel mirrored to the output bucket, so
/api/highlights can
1105         307-redirect off the app's response path (bypassing Cloud Run's ~32 MiB cap + the
ephemeral local
1106         /tmp; #463). None ? no bucket copy ? serve the local file. Uses the active storage
config for BOTH
1107         the existence check and the signing (extracted here to keep backend/main.py lean ?
P23/P24). """
1108         reel_uri = reel_object_uri(cfg, video_id, name)
1109         if not reel_uri:
1110             return None
1111         settings = _storage_settings(cfg)
1112         if storage.exists(reel_uri, config=settings):
1113             return storage.signed_url(reel_uri, config=settings, method="GET")
1114         return None
1115
1116
1117     def _execute_compile(cfg, db, job: Dict[str, Any], progress=None) -> Dict[str, Any]:
1118         """Worker handler for a ``kind='compile'`` job (#329): re-resolve the segments,
stitch the reel,
1119         and return the {status, filepath, download_url} result the SPA polls for. Returns a
1120         ``status='failed'`` result for an EXPECTED problem (missing video / no segments / a
stitch error)
1121         so the SPA always gets a clean verdict instead of a worker-crash error. """

```

```

1122     import backend.main as _main # call-time seam: VideoCompiler stays patchable on
backend.main
1123
1124     notify = progress or (lambda stage: None)
1125     params = json.loads(job.get("params") or "{}")
1126     video_id = job.get("video_id") or ""
1127     segment_ids = params.get("segment_ids") or []
1128     locale = params.get("locale")
1129
1130     notify("resolving")
1131     try:
1132         video, kept, out_name = _resolve_compile(
1133             db, cfg, video_id, segment_ids, params.get("output_filename", "")
1134         )
1135     except _CompileError as e:
1136         return {"status": "failed", "message": e.detail, "video_id": video_id}
1137
1138     stitching = getattr(cfg, "stitching", None) or StitchingConfig()
1139     output_dir = (
1140         getattr(getattr(cfg, "output", None), "output_dir", "output") or "output"
1141     )
1142     # Localize the branding watermark to the UI locale the SPA rendered in (recorded on
the job);
1143     # unknown/absent locale keeps the configured default unchanged.
1144     localized_watermark = localize_watermark(stitching.watermark, locale)
1145     compiler = _main.VideoCompiler(
1146         output_dir,
1147         begin_padding=stitching.begin_padding,
1148         end_padding=stitching.end_padding,
1149         watermark=localized_watermark,
1150     )
1151     time_pairs = [(s["start_time"], s["end_time"]) for s in kept]
1152
1153     notify("stitching")
1154     local_src = ""
1155     try:
1156         # The stored `filepath` is the SOURCE ref (gs://, s3://, gdrive://, or a local
path). Resolve to
1157         # a LOCAL file the stitcher can read ? identity for a local path, fetch-to-
scratch for a URI (the
1158         # same seam the process path uses); without it, compiling a bucket-ingested
video can't open gs://.
1159         local_src = storage.acquire_local_input(
1160             video["filepath"], getattr(cfg, "storage", None)
1161         )
1162         # The ffmpeg stitch is CPU-heavy ? serialize it so a raised drain worker count
(for remote
1163         # dispatch) doesn't run multiple compiles at once on the box (see
_LOCAL_COMPUTE_LANE).
1164         with _LOCAL_COMPUTE_LANE:
1165             out_path = compiler.compile_highlights(local_src, time_pairs, out_name)
1166     except Exception as e: # noqa: BLE001 ? a stitch/ffmpeg failure is a clean job
FAILURE, not a crash
1167         logger.error(
1168             "[compile] stitching failed for video_id=%s: %s", video_id, e, exc_info=True
1169         )
1170     return {
1171         "status": "failed",
1172         "message": f"Highlight stitching failed: {e}",
1173         "video_id": video_id,
1174     }
1175 finally:
1176     if local_src:
1177         storage.cleanup_acquired_local_input(
1178             video["filepath"], local_src, getattr(cfg, "storage", None)

```

```

1179         )
1180
1181         # #463: mirror the reel to the output bucket so /api/highlights can 307 to a signed
URL ? Cloud Run
1182         # caps a buffered response at ~32 MiB (a real reel is bigger) and the local /tmp
copy dies when the
1183         # instance recycles. A mirror-upload failure must NOT fail an otherwise-good stitch
(the endpoint
1184         # still falls back to the local file on the same instance).
1185         reel_uri = reel_object_uri(cfg, video_id, out_name)
1186         if reel_uri:
1187             try:
1188                 notify("uploading reel")
1189                 # Thread the active storage settings so the configured backend (S3
endpoint/region, GCS
1190                 # project/client) is used, not a fresh default.
1191                 storage.put(out_path, reel_uri, config=_storage_settings(cfg))
1192             except Exception as e: # noqa: BLE001 ? best-effort mirror; local stitch
already succeeded
1193                 logger.warning("[compile] could not mirror reel to %s: %s", reel_uri, e)
1194
1195         # `filepath` is server-local (not browser-reachable); also return the download URL
so the SPA needn't
1196         # hand-build it (P6 / PER_RALLY_SELECTION.md §4.3).
1197         download_url = f"/api/highlights/{quote(video_id)}/{quote(out_name)}"
1198         notify("finished")
1199         return {
1200             "status": "success",
1201             "filepath": out_path,
1202             "download_url": download_url,
1203             "video_id": video_id,
1204         }
1205

```

13. user (2026-07-08T05:05:51.816Z)

```

1      """Worker dispatcher ? pull from storage ? run the pipeline ? push outputs back.
2
3      python -m backend.worker infer --video-uri <uri> --out-uri <prefix> \
4          [--segmenter wasb_hybrid] [--config config.json] [--set k=v ...] [--scratch DIR]
5      python -m backend.worker train ...          # Gen-0 train-from-bucket (follow-up)
6
7      The URI scheme selects the storage backend via ``backend.storage`` (a bare path /
``local://``
8      is an identity passthrough; ``s3://`` / ``gcs://`` is an object store), so existing
local runs
9      are byte-for-byte unchanged. Combine with ``RALLY_DETECTOR_IMPL=stub`` and
10     ``--set indexing.skip_ai_handoff=true`` for a **secret-free, GPU-free** run (no Gemini
key) ?
11     the path validated end-to-end on DigitalOcean (``docs/SETUP_NEW_MACHINE.md``).
12
13     This is a thin orchestrator: it composes the SAME building blocks the ``backend.main``
CLI uses
14     (validator registry ? segmenter ? report), never a forked pipeline.
15     """
16
17     from __future__ import annotations
18
19     import argparse
20     import csv
21     import json
22     import logging
23     import os
24     import re
25     import sys

```



```

26     import tempfile
27     from typing import Any, List
28
29     from backend import storage
30     from backend.compute import ComputeJobSpec, get_compute_backend
31     from backend.config import (
32         StartupCheckError,
33         enforce_startup_checks,
34         load_config,
35         probe_cuda_available,
36     )
37     from backend.infrastructure.database import Database
38     from backend.pipeline.segmenters.base import segmenter_registry
39     from backend.pipeline.validators.base import registry as validator_registry
40     from backend.reporting import emit_indexer_report
41     from backend.results import Failure
42     from backend.utils.ffmpeg import ffprobe_bin # P2a: single ffmpeg/ffprobe resolver
43     from backend.utils.hashing import compute_video_id
44
45     logger = logging.getLogger("backend.worker")
46
47
48     def _coerce(v: str) -> Any:
49         """Coerce a ``--set`` string value to bool/int/float where it round-trips, else
50         str."""
51         low = v.lower()
52         if low in ("true", "false"):
53             return low == "true"
54         for cast in (int, float):
55             try:
56                 return cast(v)
57             except ValueError:
58                 pass
59         return v
60
61     def _apply_overrides(config: Any, sets: List[str]) -> None:
62         """Apply ``a.b.c=value`` overrides onto the typed config (e.g.
63         indexing.skip_ai_handoff=true)."""
64         for kv in sets or []:
65             key, sep, val = kv.partition("=")
66             if not sep:
67                 raise ValueError(f"--set expects k=v, got {kv!r}")
68             parts = key.split(".")
69             obj = config
70             for p in parts[:-1]:
71                 obj = getattr(obj, p)
72             setattr(obj, parts[-1], _coerce(val))
73
74     def _write_intervals_csv(segments: List[dict], path: str) -> None:
75         """Decimal-second ``[start,end)`` rallies + shot count / ending reason ? the join-
76         friendly
77         artifact (same schema as the rally-annotator labels)."""
78         with open(path, "w", newline="") as f:
79             w = csv.writer(f)
80             w.writerow(["start", "end", "shot_count", "ending_reason"])
81             for s in segments:
82                 w.writerow(
83                     [
84                         f"{float(s.get('start_time', 0.0)):.3f}",
85                         f"{float(s.get('end_time', 0.0)):.3f}",
86                         s.get("shot_count", ""),
87                         s.get("ending_reason", s.get("shot_count_confidence", "")),
88                     ]

```

```

88         )
89
90
91     def _write_report_json(
92         video_id: str, segments: List[dict], status: str, path: str, video_uri: str = ""
93     ) -> None:
94         with open(path, "w") as f:
95             json.dump(
96                 {
97                     "video_id": video_id,
98                     "status": status,
99                     "segment_count": len(segments),
100                     # The SOURCE the detection ran on (CP-rebuild P3, #485): a restart-wiped
101                     # control plane re-hydrating from outputs/<md5>/ restores the video row
102                     # with its truthful path ? without this, only the segments recover.
103                     "video_uri": video_uri,
104                     "segments": [
105                         {
106                             "start": float(s.get("start_time", 0.0)),
107                             "end": float(s.get("end_time", 0.0)),
108                         }
109                         for s in segments
110                     ],
111                 },
112                 f,
113                 indent=2,
114             )
115
116
117     def _serialize_and_push_outputs(
118         scratch: str,
119         out_uri: str,
120         video_id: str,
121         segments: List[dict],
122         status: str,
123         storage_cfg: dict,
124         video_uri: str = "",
125     ) -> List[str]:
126         """Write intervals.csv + report.json under ``<scratch>/outputs/<video_id>/`` and
127         push both to
128         ``<out_uri>/outputs/<video_id>/``; return the pushed URIs.
129
130         Shared by the SUCCESS path and the failure path (``_fail``) so a detector
131         timeout/abort leaves a
132         ``status="failed"`` report.json in the bucket ? the control plane's source of truth
133         (``orchestration.cached_process_result`` / ``probe_process_cache`` treat any
134         non-``success``
135         report as "no usable result" ? retry, and the restart-recovery poll waits for
136         report.json to
137         EXIST) ? instead of a missing artifact or a false 0-segment ``"success"``.
138         """
139         out_local = os.path.join(scratch, "outputs", video_id)
140         os.makedirs(out_local, exist_ok=True)
141         _write_intervals_csv(segments, os.path.join(out_local, "intervals.csv"))
142         _write_report_json(
143             video_id,
144             segments,
145             status,
146             os.path.join(out_local, "report.json"),
147             video_uri=video_uri,
148         )
149         out_prefix = out_uri.rstrip("/")
150         pushed: List[str] = []
151         for fn in sorted(os.listdir(out_local)):
152             uri = f"{out_prefix}/outputs/{video_id}/{fn}"
153             storage.put(os.path.join(out_local, fn), uri, config=storage_cfg)

```

```

149         pushed.append(uri)
150     return pushed
151
152
153     def _run_startup_checks(config: Any) -> bool:
154         """M5 per-profile startup checks for the worker (the compute process, so it probes
CUDA
155         best-effort). Logs warnings; returns False on a FATAL coherence error so the caller
aborts
156         cleanly (rc 2) rather than run an expensive job under an incoherent profile. A true
no-op
157         (returns True, ZERO work) when no deployment.profile is selected.
158
159         The CUDA probe is gated behind an active profile: ``probe_cuda_available()`` lazily
imports
160         torch, and pre-M5 the worker's infer/train path was torch-free (detection is
delegated to a
161         WSL/native subprocess). Probing only when a profile is set keeps the default-OFF
path a true
162         no-op of work, not just of output."""
163         profile = getattr(getattr(config, "deployment", None), "profile", "") or ""
164         cuda = (
165             probe_cuda_available() if profile else None
166         ) # torch-free on the default-OFF path
167         try:
168             enforce_startup_checks(config, cuda_available=cuda, logger=logger)
169             return True
170         except StartupCheckError:
171             return False # already logged at ERROR by enforce_startup_checks
172
173
174     def _dispatch_remote(compute: Any, args: argparse.Namespace) -> int:
175         """M6: hand ONE infer job to a remote ComputeBackend (Cloud Run) and return its rc
once the
176         outputs are durably in the bucket (the backend bucket-polls the done-marker). The
dispatched
177         container runs this SAME cmd_infer INLINE (compute_target forced local_gpu) ? never
re-dispatch.
178
179         A remote job that never writes a completion marker exits via one of two CLEAN codes
(distinct
180         from the local 2/3), never a raw traceback:
181         * rc 4 ? the bucket poll exhausted its budget (``PolicyTimeout``): the execution
may still be
182         running (the shim best-effort cancels it).
183         * rc 5 ? the execution has TERMINATED without a marker
(``RemoteExecutionFailed``): a worker
184         crash/OOM/pre-marker abort caught FAST by the Executions-API probe, so instead
of the
185         ~65-min opaque wait the operator gets an immediate, log-pointing failure."""
186         from backend.compute import RemoteExecutionFailed
187         from backend.infrastructure.execution_policy import PolicyTimeout
188
189         spec = ComputeJobSpec(
190             video_uri=args.video_uri,
191             out_uri=args.out_uri,
192             segmenter=args.segmenter,
193             config_overrides=tuple(args.set or ()),
194             skip_validation=bool(getattr(args, "skip_validation", False)),
195         )
196         try:
197             result = compute.run_infer(spec)
198         except RemoteExecutionFailed as e:
199             logger.error(
200                 "[worker] remote execution terminated without a completion marker: %s", e

```

```

201         )
202         return 5
203     except PolicyTimeout as e:
204         logger.warning(
205             "[worker] remote compute did not complete (no marker within budget): %s", e
206         )
207         return 4
208     logger.info(
209         "[worker] remote compute done: video_id=%s rc=%d outputs=%s",
210         result.video_id,
211         result.returncode,
212         result.outputs_uri,
213     )
214     return result.returncode
215
216
217     def _write_done_marker(
218         done_uri: str,
219         video_id: str,
220         returncode: int,
221         segment_count: int,
222         status: str,
223         storage_cfg: dict,
224         scratch: str,
225     ) -> None:
226         """M6: write a tiny JSON completion marker to ``done_uri`` (via the storage layer)
227         for a remote
228         dispatcher's bucket-poll. Best-effort ? never raises into the worker's success
229         return."""
230         try:
231             marker = {
232                 "video_id": video_id,
233                 "returncode": returncode,
234                 "segment_count": segment_count,
235                 "status": status,
236             }
237             local = os.path.join(scratch, "_done.json")
238             with open(local, "w", encoding="utf-8") as f:
239                 json.dump(marker, f)
240             storage.put(local, done_uri, config=storage_cfg)
241             logger.info("[worker] wrote completion marker -> %s", done_uri)
242         except Exception as e: # never fail a good run because the marker push hiccuped
243             logger.warning("[worker] could not write done-marker %s: %s", done_uri, e)
244
245
246     def cmd_infer(args: argparse.Namespace) -> int:
247         config = load_config(args.config) if args.config else load_config()
248         _apply_overrides(config, args.set)
249         if not _run_startup_checks(config):
250             return 2
251
252         # M6: a cloud compute_target DISPATCHES to a Cloud Run Job instead of running in-
253         # process.
254         # DEFAULT-OFF: get_compute_backend returns None for ""/local_gpu/cpu_mock ? fall
255         # through to the
256         # unchanged inline path below. The dispatched container runs THIS cmd_infer inline
257         # (the
258         # dispatcher forces compute_target=local_gpu), so it never recursively re-
259         # dispatches.
260         compute = get_compute_backend(config)
261         if compute is not None:
262             return _dispatch_remote(compute, args)
263
264         storage_cfg = config.storage.model_dump()

```

```

260     scratch = args.scratch or tempfile.mkdtemp(prefix="rally-worker-")
261     os.makedirs(scratch, exist_ok=True)
262     config.output.output_dir = scratch
263
264     # 1. PULL ? local:// / bare path = identity passthrough; s3:// / gcs:// = download
to scratch.
265     local_video = storage.fetch(
266         args.video_uri, os.path.join(scratch, "in.mp4"), config=storage_cfg
267     )
268     print(f"[worker] pulled {args.video_uri} -> {local_video}")
269
270     # 2. IDENTITY ? whole-file MD5 (the labels?video join key; pulled bytes must be
byte-identical).
271     video_id = compute_video_id(local_video)
272
273     # 3. VALIDATE (same registry as the main CLI) ? UNLESS the dispatcher already did.
274     # On the cloud-serving path the CONTROL-PLANE is the authoritative validation gate:
it runs the
275     # SAME validators with ITS resolved config and only dispatches on PASS. Re-
validating here is
276     # redundant, decodes the whole video a SECOND time, and is outright WRONG when the
worker's config
277     # resolves different thresholds than the control-plane's ? observed 2026-07-07: a
5312x2988 clip
278     # (sharpness 11.9) PASSED the control-plane at validation.min_blur_threshold=8.0,
then the worker
279     # REJECTED it at its own default 20.0, so a SUCCESSFUL GPU dispatch failed after the
fact.
280     # --skip-validation (set by the dispatcher in orchestration._maybe_dispatch_remote)
skips it; the
281     # direct ``worker infer`` CLI has no control-plane, so it still validates by
default.
282     if getattr(args, "skip_validation", False):
283         print(
284             "[worker] validation SKIPPED (--skip-validation): the dispatcher is the
authoritative "
285             "gate and already validated this input."
286         )
287         val_results, val_context = [], {}
288     else:
289         val_results, val_context = validator_registry.run_all_with_context(
290             local_video, config
291         )
292     failures = [r for r in val_results if not r.passed]
293     db = Database(os.path.join(scratch, config.output.db_name))
294     db.add_video(
295         video_id,
296         local_video,
297         "failed" if failures else "processed",
298         [r.model_dump() for r in val_results],
299     )
300
301     def _fail(rc: int) -> int:
302         # A worker-side FAILURE (validation, unknown segmenter, or a detector
timeout/abort the
303         # segmenter surfaced as a Failure) must NOT masquerade as a 0-rally "success".
Write + push
304         # a status="failed" report.json/intervals.csv (0 segments) so the control
plane's bucket
305         # truth (submit-cache probe, restart-recovery) and the alpha-user result surface
see an
306         # explicit FAILURE ? the wf294 silent-failure fix. Best-effort: a push hiccup
must never
307         # mask the original failure. THEN the M6 done-marker (real rc + status) so a
remote

```

```

308         # dispatcher's bucket-poll terminates FAST instead of hanging until its budget
expires.
309     try:
310         _serialize_and_push_outputs(
311             scratch,
312             args.out_uri,
313             video_id,
314             [],
315             "failed",
316             storage_cfg,
317             str(getattr(args, "video_uri", "") or ""),
318         )
319     except Exception as e: # never mask the real failure on an artifact-push hiccup
320         logger.warning("[worker] could not push failure artifacts: %s", e)
321     # DEFAULT-OFF: no --done-uri ? no marker written (the unchanged local run).
322     if getattr(args, "done_uri", None):
323         _write_done_marker(
324             args.done_uri, video_id, rc, 0, "failed", storage_cfg, scratch
325         )
326     return rc
327
328     segments: List[dict] = []
329     segmenter_actual = None
330     segmentation_ran = False
331     status = "failed"
332     try:
333         if failures:
334             print(
335                 "[worker] validation FAILED: "
336                 + "; ".join(f"{r.validator_name}: {r.message}" for r in failures)
337             )
338             return _fail(2)
339         seg_name = args.segmenter or config.indexing.default_segmenter
340         segmenter_cls = segmenter_registry.get(seg_name)
341         if not segmenter_cls:
342             print(
343                 f"[worker] unknown segmenter: {seg_name!r} "
344                 f"(have {list(segmenter_registry.list_available())})"
345             )
346             return _fail(2)
347         segmenter_actual = seg_name
348         result = segmenter_cls(db, config).process_video(
349             video_id, local_video, max_frames=args.max_frames
350         )
351         segmentation_ran = True
352         if isinstance(result, Failure):
353             print(f"[worker] segmenter FAILED: {result.message}")
354             return _fail(3)
355         segments = result.value if hasattr(result, "value") else result
356         status = "success"
357         # Loud, never-silent diagnostic: a 0-segment success is almost always a
misconfig, not an
358         # empty match ? the cold-start failure mode (the substrate detector yielded no
usable
359         # trajectories). Explain it so a run is analysable from the log alone (re-run -v
for more).
360         if not segments:
361             detector_impl = (
362                 os.environ.get("RALLY_DETECTOR_IMPL")
363                 or getattr(config.indexing, "detector_impl", None)
364                 or "auto"
365             )
366             logger.warning(
367                 "infer produced 0 segments for video_id=%s (segmenter=%s,
detector_impl=%s). "

```

```

368         "The detector substrate yielded no usable trajectories/windows ? check
the detector "
369         "logs above; common causes: a missing ${AI_PROVIDER}_API_KEY with "
370         "skip_ai_handoff=False (the segmenter bails BEFORE detection ? set the
key or "
371         "indexing.skip_ai_handoff=true; cloud-serving now fails this at
startup), native "
372         "runner UNAVAILABLE (weights/repo/WASB_PYTHON env), the WASB env
produced no "
373         "detections, or the input resolution differs from the tuned proxy
resolution. "
374         "Re-run with -v for the native command + frame counts.",
375         video_id,
376         segmenter_actual,
377         detector_impl,
378     )
379     finally:
380         # Best-effort per-video observability report (never raises) ? parity with the
main CLI.
381         emit_indexer_report(
382             db=db,
383             video_id=video_id,
384             video_path=local_video,
385             config=config,
386             val_results=val_results,
387             val_context=val_context,
388             segmenter_requested=args.segmenter,
389             segmenter_actual=segmenter_actual,
390             was_reingest=False,
391             segmentation_ran=segmentation_ran,
392         )
393
394     # 5. SERIALIZE + 6. PUSH (outputs/<video_id>/?). idempotent on video_id.
395     pushed = _serialize_and_push_outputs(
396         scratch,
397         args.out_uri,
398         video_id,
399         segments,
400         status,
401         storage_cfg,
402         str(getattr(args, "video_uri", "") or ""),
403     )
404
405     print(
406         f"[worker] infer done: {len(segments)} segment(s); pushed {len(pushed)}
artifact(s):"
407     )
408     for u in pushed:
409         print("    ", u)
410
411     # M6: a remote dispatcher passes --done-uri; write the SUCCESS completion marker so
its bucket-poll
412     # terminates (carries video_id ? the dispatcher's output path + rc=0). Worker
FAILURES (rc 2/3)
413     # write their own marker via _fail() above, so the dispatcher fails FAST either way
(no 30-min hang
414     # on a bad video). DEFAULT-OFF: no --done-uri ? no marker written (the unchanged
local run).
415     if getattr(args, "done_uri", None):
416         _write_done_marker(
417             args.done_uri, video_id, 0, len(segments), status, storage_cfg, scratch
418         )
419     return 0
420
421

```

```

422     #: Video extensions accepted under <corpus>/videos/ (mirrors the labels/ '.csv' filter
so a
423     #: per-video sidecar ? .srt/.json/thumbnail ? can't shadow the real video on a stem
collision).
424     _VIDEO_EXTS = (".mp4", ".mov", ".mkv", ".avi", ".webm", ".m4v")
425
426     #: Leading ISO-date prefix on a canonical label name (`labels/<YYYY-MM-
DD>_<name>.rallies.csv`,
427     #: docs/DATA_IN_GCS.md §7.1). It disambiguates recycled GoPro names by authoring date
but is NOT
428     #: part of the clip's join key, so it's stripped when deriving the videos/ stem.
429     _LABEL_DATE_PREFIX = re.compile(r"^\d{4}-\d{2}-\d{2}_")
430
431
432     def _label_clip_name(fname: str) -> str:
433         """Derive the clip name (the `videos/<name>.<ext>` join key) from a label
filename.
434
435         Strips both the `.rallies.csv` / `.csv` suffix and an optional leading ISO-date
prefix, so a
436         canonical dated label and a legacy bare label resolve to the SAME bare clip stem:
437
438         * `2026-06-22_GX020094.rallies.csv` -> `GX020094` (canonical, date-
stamped)
439         * `GX020094.rallies.csv` -> `GX020094` (legacy bare ? same
stem)
440         * `GX020094_proxy.rallies.csv` -> `GX020094_proxy` (`_proxy` is part of
the name, kept)
441         * `2026-06-21_mahadevpura_1.rallies.csv` -> `mahadevpura_1` (underscores in
<name> survive)
442
443         Without the date strip, `--trajectory-source detector` can't match a dated label
to its bare
444         video stem, so every clip is skipped and train aborts with `<2 videos`
(DATA_IN_GCS §7b #1).
445         """
446         name = fname.replace(".rallies.csv", "").replace(".csv", "")
447         return _LABEL_DATE_PREFIX.sub("", name)
448
449
450     def _probe_video_fps_width(path: str):
451         """Robust ffprobe (fps, frame_width) for a real video. Returns None on ANY failure.
452
453         The detector trajectory is frame-indexed; the harness maps frame->time via fps, so a
WRONG
454         fps silently mis-aligns every rally label (e.g. a 60fps clip read as 30 doubles
every time).
455         cv2's `_probe_fps_width` silently defaults to (30, 1280) on a read miss ? fine for
the
456         inference report (it flags the fallback) but corrupting for REAL training data. So
here we
457         probe with ffprobe (same tool as `get_video_duration`) and return None so the
caller SKIPS
458         rather than guesses.
459         """
460         import json
461         import subprocess
462
463         try:
464             out = subprocess.check_output(
465                 [
466                     ffprobe_bin(),
467                     "-v",
468                     "error",
469                     "-select_streams",

```



```

470         "v:0",
471         "-show_entries",
472         "stream=r_frame_rate,width",
473         "-of",
474         "json",
475         path,
476     ],
477     stderr=subprocess.DEVNULL,
478 ).decode()
479 st = (json.loads(out).get("streams") or [{}])[0]
480 num, _, den = str(st.get("r_frame_rate", "0/1")).partition("/")
481 fps = float(num) / float(den or "1")
482 width = float(st.get("width") or 0)
483 if fps > 0 and width > 0:
484     return fps, width
485 except (
486     subprocess.SubprocessError,
487     ValueError,
488     KeyError,
489     OSError,
490     json.JSONDecodeError,
491 ):
492     pass
493 return None
494
495
496 def _resolve_train_detector(args: argparse.Namespace, config: Any):
497     """Build the trajectory DetectorRunner for `--trajectory-source detector`, honouring
498     detector_impl (RALLY_DETECTOR_IMPL / indexing.detector_impl: auto=WSL, native=GPU
499     box,
500     stub=CPU mock). Returns the runner, or prints an error + returns None.
501
502     Reuses the segmenter's `_resolve_runner` seam so the worker and the live CLI build
503     the
504     detector identically (one source of truth). A non-trajectory segmenter (e.g. yolo)
505     has
506     no shuttle detector and is rejected.
507     """
508     from backend.pipeline.segmenters.trajectory_hybrid import TrajectoryHybridSegmenter
509
510     seg_cls = segmenter_registry.get(args.segmenter)
511     if not seg_cls:
512         print(
513             f"[worker] unknown segmenter: {args.segmenter!r} "
514             f"(have {list(segmenter_registry.list_available())})"
515         )
516         return None
517     seg = seg_cls(None, config)
518     if not isinstance(seg, TrajectoryHybridSegmenter):
519         print(
520             f"[worker] --segmenter {args.segmenter!r} has no shuttle detector; "
521             f"use a trajectory hybrid (wasb_hybrid / tracknet_hybrid / fusion_hybrid)."
522         )
523         return None
524     try:
525         runner, _ = seg._resolve_runner(config.indexing)
526     except NotImplementedError as e:
527         # e.g. detector_impl='native' for a family without a native runner (tracknet, or
528         # fusion with a tracknet substrate). Fail cleanly (rc!=0), not with a raw
529         traceback.
530         print(f"[worker] detector not available: {e}")
531         return None
532     ok, msg = runner.healthcheck()
533     if not ok:
534         print(f"[worker] detector environment not ready: {msg}")

```

```

531         return None
532     print(f"[worker] detector ready ({args.segmenter}): {msg}")
533     return runner
534
535
536 def cmd_train(args: argparse.Namespace) -> int:
537     """Gen-0 train-from-bucket: pull labels (+ videos) -> trajectories -> manifest ->
shell out
538     to training.gen0.harness (backend must NOT import training) -> push the artifact
bundle.
539
540     Two trajectory sources (the train pipeline + harness are identical either way):
541     * ``synth`` (default) ? the CPU "mock GPU": deterministic, label-consistent
synthetic
542         shuttle paths (no GPU/WSL), so the whole pipeline is validated on a CPU box /
CI.
543     * ``detector`` ? REAL shuttle trajectories: pull each video and run the resolved
544         DetectorRunner (native WASB on a GPU box, or stub for a GPU-free wiring test).
This
545         is the M3.5 "first real training" path; only the trajectory source flips.
546     """
547     import subprocess
548
549     config = load_config(args.config) if args.config else load_config()
550     _apply_overrides(config, args.set)
551     if not _run_startup_checks(config):
552         return 2
553     storage_cfg = config.storage.model_dump()
554
555     scratch = args.scratch or tempfile.mkdtemp(prefix="rally-train-")
556     labels_dir = os.path.join(scratch, "labels")
557     traj_dir = os.path.join(scratch, "trajectories")
558     videos_dir = os.path.join(scratch, "videos")
559     os.makedirs(labels_dir, exist_ok=True)
560     os.makedirs(traj_dir, exist_ok=True)
561
562     corpus = args.corpus_uri.rstrip("/")
563     label_uris = sorted(
564         u
565         for u in storage.list_uris(f"{corpus}/labels/", config=storage_cfg)
566         if u.lower().endswith(".csv")
567     )
568     if len(label_uris) < 2:
569         print(
570             f"[worker] need >=2 labeled videos for leave-one-video-out; found
{len(label_uris)} "
571             f"under {corpus}/labels/"
572         )
573         return 2
574
575     # Real-detector source: resolve the runner once + index the bucket's videos/ by
stem.
576     runner = None
577     video_by_stem: dict = {}
578     if args.trajectory_source == "detector":
579         os.makedirs(videos_dir, exist_ok=True)
580         runner = _resolve_train_detector(args, config)
581         if runner is None:
582             return 2
583         for vuri in storage.list_uris(f"{corpus}/videos/", config=storage_cfg):
584             vname = storage.parse_uri(vuri).name
585             if not vname.lower().endswith(_VIDEO_EXTS):
586                 continue # skip sidecars (.srt/.json/thumbnails) so they can't shadow
the video
587             video_by_stem[os.path.splitext(vname)[0]] = vuri

```

```

588
589 print(f"[worker] trajectory source: {args.trajectory_source}")
590 specs: List[dict] = []
591 for uri in label_uris:
592     fname = storage.parse_uri(uri).name # e.g. 2026-06-22_GX020094.rallies.csv
593     name = _label_clip_name(fname) # -> GX020094 (date prefix stripped)
594     local_label = storage.fetch(
595         uri, os.path.join(labels_dir, fname), config=storage_cfg
596     )
597     if args.trajectory_source == "detector":
598         vuri = video_by_stem.get(name)
599         if not vuri:
600             print(
601                 f"[worker] no video for {name!r} under {corpus}/videos/ ? skipping."
602             )
603             continue
604         local_video = storage.fetch(
605             vuri,
606             os.path.join(videos_dir, storage.parse_uri(vuri).name),
607             config=storage_cfg,
608         )
609         video_id = compute_video_id(local_video)
610         # Real fps/width drive the trajectory frame->time mapping; PROBE (don't
guess) ? a
611         # wrong fps silently mis-aligns every label, so skip the video if ffprobe
can't read it.
612         meta = _probe_video_fps_width(local_video)
613         if meta is None:
614             print(
615                 f"[worker] could not probe fps/width for {name!r} (ffprobe) ?
skipping (won't guess)."
616             )
617             continue
618         fps, frame_width = meta
619         traj = runner.run_predict(local_video, traj_dir, video_id=video_id)
620         if not traj:
621             print(
622                 f"[worker] detector produced no trajectory for {name!r} ? skipping."
623             )
624             continue
625         specs.append(
626             {
627                 "name": name,
628                 "trajectory": traj,
629                 "labels": local_label,
630                 "fps": fps,
631                 "frame_width": frame_width,
632             }
633         )
634     else:
635         from backend.pipeline.detectors.stub_runner import (
636             synth_trajectory_from_labels,
637         )
638
639         traj = os.path.join(traj_dir, f"{name}_traj.csv")
640         synth_trajectory_from_labels(
641             local_label, traj, fps=args.fps, frame_width=args.frame_width
642         )
643         specs.append(
644             {
645                 "name": name,
646                 "trajectory": traj,
647                 "labels": local_label,
648                 "fps": args.fps,
649                 "frame_width": args.frame_width,

```

```

650         }
651     )
652
653     if len(specs) < 2:
654         print(
655             f"[worker] need >=2 videos with trajectories for leave-one-video-out; got
{len(specs)}."
656         )
657         return 2
658     manifest_path = os.path.join(scratch, "manifest.json")
659     with open(manifest_path, "w", encoding="utf-8") as f:
660         json.dump(specs, f, indent=2)
661     src_label = (
662         "real detector" if args.trajectory_source == "detector" else "CPU-mock stub"
663     )
664     print(f"[worker] built manifest: {len(specs)} videos ({src_label} trajectories)")
665
666     out_dir = os.path.join(scratch, "artifacts")
667     cmd = [
668         sys.executable,
669         "-m",
670         "training.gen0.harness",
671         "--manifest",
672         manifest_path,
673         "--out",
674         out_dir,
675         "--version",
676         args.version,
677     ]
678     if args.floor:
679         floor_local = (
680             storage.fetch(
681                 args.floor, os.path.join(scratch, "floor.json"), config=storage_cfg
682             )
683             if "://" in args.floor
684             else args.floor
685         )
686         cmd += ["--floor", floor_local]
687     print("[worker] training:", " ".join(cmd))
688     rc = subprocess.run(cmd).returncode
689     if rc != 0:
690         print(f"[worker] gen0 harness failed (rc={rc})")
691         return rc
692
693     # PUSH the versioned artifact bundle (weights.json + manifest.json) to the bucket.
694     bundle = os.path.join(out_dir, args.version)
695     out_prefix = args.out_uri.rstrip("/")
696     pushed = []
697     for fn in sorted(os.listdir(bundle)):
698         uri = f"{out_prefix}/weights/{args.version}/{fn}"
699         storage.put(os.path.join(bundle, fn), uri, config=storage_cfg)
700         pushed.append(uri)
701     print(f"[worker] train done: pushed {len(pushed)} artifact(s):")
702     for u in pushed:
703         print("    ", u)
704     return 0
705
706
707 def build_parser() -> argparse.ArgumentParser:
708     p = argparse.ArgumentParser(
709         prog="python -m backend.worker",
710         description="Storage-aware worker: pull ? run pipeline ? push.",
711     )
712     p.add_argument(
713         "-v",

```

```

714         "--verbose",
715         action="store_true",
716         help="DEBUG logging (the native detector command, frame counts, per-stage
detail)",
717     )
718     sub = p.add_subparsers(dest="cmd", required=True)
719
720     pi = sub.add_parser("infer", help="run inference on one video URI and push outputs")
721     pi.add_argument(
722         "--video-uri", required=True, help="local path / local:// / s3:// / gcs://"
723     )
724     pi.add_argument(
725         "--out-uri",
726         required=True,
727         help="output prefix (outputs/<video_id>/? is appended)",
728     )
729     pi.add_argument(
730         "--segmenter", default=None, help="default: indexing.default_segmenter"
731     )
732     pi.add_argument(
733         "--config", default=None, help="config.json path (default: ./config.json)"
734     )
735     pi.add_argument(
736         "--scratch", default=None, help="local scratch dir (default: a temp dir)"
737     )
738     pi.add_argument("--max-frames", type=int, default=None)
739     pi.add_argument(
740         "--done-uri",
741         default=None,
742         help="M6: storage URI for a completion MARKER (video_id + rc) written on a "
743         "successful run, so a remote (Cloud Run) dispatcher's bucket-poll terminates. "
744         "Default-OFF: omit it for the normal local run.",
745     )
746     pi.add_argument(
747         "--set",
748         action="append",
749         default=[],
750         metavar="k=v",
751         help="config override, e.g. --set indexing.skip_ai_handoff=true",
752     )
753     pi.add_argument(
754         "--skip-validation",
755         action="store_true",
756         help="skip the worker's re-validation: the dispatcher (control-plane) already
ran the "
757         "validators with its resolved config and only dispatches on PASS, so re-
validating here is "
758         "redundant, decodes the video again, and can contradict the gate. Default-OFF:
the direct "
759         "`worker infer` CLI still validates.",
760     )
761     pi.set_defaults(func=cmd_infer)
762
763     pt = sub.add_parser(
764         "train",
765         help="Gen-0 train-from-bucket (labels [+videos] -> trajectories -> harness)",
766     )
767     pt.add_argument(
768         "--corpus-uri",
769         required=True,
770         help="prefix containing labels/ (+ videos/ for "
771         "--trajectory-source detector), e.g. s3://bucket or gcs://bucket",
772     )
773     pt.add_argument(
774         "--out-uri",

```

```

775         required=True,
776         help="output prefix (weights/<version>/ is appended)",
777     )
778     pt.add_argument("--version", required=True, help="semver for the artifact bundle")
779     pt.add_argument(
780         "--trajectory-source",
781         choices=["synth", "detector"],
782         default="synth",
783         help="synth = CPU-mock label-consistent trajectories (default, no GPU); "
784         "detector = REAL shuttle trajectories via the resolved DetectorRunner "
785         "(needs videos/ in the corpus; honours detector_impl).",
786     )
787     pt.add_argument(
788         "--segmenter",
789         default="wasb_hybrid",
790         help="trajectory hybrid whose detector is used for --trajectory-source",
791     )
792     pt.add_argument(
793         "--floor", default=None, help="floor baseline JSON (local path or storage URI)"
794     )
795     pt.add_argument("--config", default=None)
796     pt.add_argument("--scratch", default=None)
797     pt.add_argument(
798         "--fps",
799         type=float,
800         default=30.0,
801         help="synth trajectory fps (detector probes the video)",
802     )
803     pt.add_argument(
804         "--frame-width",
805         type=float,
806         default=1280.0,
807         help="synth trajectory frame width (detector probes the video)",
808     )
809     pt.add_argument("--set", action="append", default=[], metavar="k=v")
810     pt.set_defaults(func=cmd_train)
811     return p
812
813
814     def _setup_logging(verbose: bool) -> None:
815         """Configure logging for the worker CLI so the pipeline's existing INFO/ERROR
816         diagnostics are
817         VISIBLE ? the worker previously ran with no logging config, so every
818         segmenter/detector log
819         (detector choice, 'Running native WASB inference', availability failures) was
820         silently dropped.
821         That is exactly why a 0-rally cold-start gave nothing to analyse. Configured here
822         (the CLI
823         entry), not at import, so importing the module never reconfigures a caller's
824         logging."""
825         # force=True: basicConfig is a no-op if the root already has handlers (e.g. a
826         library
827         # configured logging at import), which would leave the level at WARNING and re-
828         suppress the
829         # diagnostics. Forcing reconfiguration means the worker reliably controls its own
830         verbosity.
831         logging.basicConfig(
832             level=logging.DEBUG if verbose else logging.INFO,
833             format="%(asctime)s [%(levelname)s] %(name)s: %(message)s",
834             datefmt="%H:%M:%S",
835             force=True,
836         )

```

```

831     def main(argv=None) -> int:
832         args = build_parser().parse_args(argv)
833         _setup_logging(getattr(args, "verbose", False))
834         return args.func(args)
835
836
837     if __name__ == "__main__":
838         sys.exit(main())
839

```

14. assistant (2026-07-08T05:05:58.441Z)

This is a large task. Let me set up a todo list and continue studying the key files: the compute backend, config models, VideoCompiler, and the job dispatch layer.

15. user (2026-07-08T05:05:59.457Z)

```

1     """`CloudRunCompute` ? dispatch one inference job to a Cloud Run GPU **Job** and
bucket-poll
2     for its completion (COMPUTE_DECOUPLED_SERVING M6, decisions D16 + D7).
3
4     Flow (the cloud-serving compute path):
5     1. **Dispatch** a Cloud Run Job execution whose container runs ``python -m
backend.worker infer
6         --video-uri <gs> --out-uri <gs> --done-uri <gs> ?`` (the SAME worker, on the L4
GPU).
7     2. **Bucket-poll** ``done_uri`` via ``execution_policy.poll_until`` + the storage
layer (D16 ? the
8         completion signal is a tiny marker object the worker writes, NOT the Cloud Run API
status, so
9         it is storage-backed and reuses the existing park/resume primitives). Scale-to-zero
(D7): the
10        cold-start is absorbed by the async poll. On a poll TIMEOUT (``max_wait_sec`` can
undercut the
11        job's own ``--task-timeout`` budget) the shim does ONE final marker check ? a slow-
but-healthy
12        worker that finished just past the deadline is rescued, not wasted ? then best-
effort CANCELS
13        the still-running execution so the L4 doesn't keep billing until ``--task-
timeout``.
14        *Fast negative path:* while the marker stays the SOLE success signal, each poll
also probes the
15        Cloud Run Executions API ? if the execution has TERMINATED without ever writing a
marker (a
16        worker crash/OOM/argparse-abort before the marker push), the shim fails fast +
diagnostically
17        (``RemoteExecutionFailed``) instead of spinning the full ~65-min budget on a futile
wait.
18    3. **Read** the marker (it carries ``video_id`` + the worker's returncode) ? fetch
19        ``outputs/<video_id>/report.json`` ? ``ComputeResult``.
20
21    The dispatched container is forced to run **INLINE** (``compute_target=local_gpu`` +
native detector)
22    so it never recursively re-dispatches ? the dispatcher decides cloud, the container just
computes.
23
24    The Cloud Run trigger is an **injected client** (``CloudRunJobClient``) so the whole
shim is unit-tested
25    with a fake (no GCP). The real ``GoogleCloudRunJobClient`` lazy-imports ``google-cloud-
run`` and is
26    owner-verified on the M7 real run.
27    """
28
29    from __future__ import annotations
30

```

```

31     import json
32     import logging
33     import os
34     import tempfile
35     import uuid
36     from abc import ABC, abstractmethod
37     from typing import Any, Callable, List, Optional
38
39     from backend.compute.base import (
40         ComputeBackend,
41         ComputeJobSpec,
42         ComputeResult,
43         RemoteExecutionFailed,
44     )
45     from backend.infrastructure.execution_policy import (
46         DEFAULT_POLICY,
47         ExecutionPolicy,
48         PolicyTimeout,
49         poll_until,
50     )
51     from backend.storage import fetch, get_storage, parse_uri
52
53     LOG = logging.getLogger("compute.cloud_run")
54
55     # The container worker must run INLINE on the GPU ? never re-enter the cloud-dispatch
56     # branch.
57     # (compute_target=local_gpu disarms the cmd_infer dispatch guard; native = the L4 WASB
58     # runner.)
59     _DEFAULT_CONTAINER_OVERRIDES: tuple[str, ...] = (
60         "deployment.compute_target=local_gpu",
61         "indexing.detector_impl=native",
62     )
63
64     class CloudRunJobClient(ABC):
65         """Triggers a Cloud Run **Job** execution with per-execution container arg
66         overrides.
67
68         Abstracted so ``CloudRunCompute`` is testable with a fake (no GCP). A real
69         implementation
70         overrides the job's container ``args`` for this execution and starts it."""
71
72         @abstractmethod
73         def run_job(
74             self,
75             job_name: str,
76             argv: List[str],
77             *,
78             region: str,
79             project: Optional[str] = None,
80         ) -> str:
81             """Start a Cloud Run Job execution running ``python -m backend.worker <argv>``;
82             return an
83             execution id/name (for logs). Does NOT block on completion (the caller bucket-
84             polls)."""
85
86         @abstractmethod
87         def cancel_execution(self, execution: str) -> None:
88             """Cancel a running execution by the name ``run_job`` returned. Called when the
89             bucket-poll
90             times out so the abandoned execution doesn't keep billing the GPU until
91             ``--task-timeout``.
92             May raise ? the caller treats every failure as best-effort (the original poll
93             timeout is
94             NEVER masked by a cancel error)."""

```



```

87
88     def execution_terminal_state(self, execution: str) -> Optional[str]:
89         """Return the execution's TERMINAL state as a short label (``Failed`` /
90         ``Cancelled`` /
91         ``Succeeded``) if it has finished, else ``None`` (still pending/running, or
the state
92         can't be determined).
93
94         ADVISORY, not authoritative: it feeds ``run_infer``'s fast negative path so a
worker that
95         terminated WITHOUT writing the done-marker (crash/OOM/pre-marker abort) fails
fast instead
96         of bucket-polling the full budget. The done-marker stays the ONLY success
signal.
97
98         NON-abstract with a default of ``None`` (additive / default-OFF): a client that
can't ? or
99         chooses not to ? poll the Executions API degrades transparently to today's
marker-only
100        polling. Implementations MAY raise on a transient API error ? ``run_infer``
treats a raise
101        as ``None`` (unknown ? keep waiting), so a flaky status API never fails a
healthy run."""
102        return None
103
104    class GoogleCloudRunJobClient(CloudRunJobClient):
105        """Real client ? lazy-imports ``google-cloud-run``. Owner-verified on the M7 real
run; CI uses
106        a fake, so this code path is never exercised in tests (the SDK isn't a test
dependency)."""
107
108        def run_job(
109            self,
110            job_name: str,
111            argv: List[str],
112            *,
113            region: str,
114            project: Optional[str] = None,
115        ) -> str:
116            # A None/empty project would build a bogus "projects/None/locations/..." path
(job_path
117            # requires a real str). Fail loudly BEFORE the SDK import so the misconfig is
exercisable
118            # in CI (where google-cloud-run is absent) and so mypy narrows project to str
for job_path.
119            if not project:
120                raise RuntimeError(
121                    "GOOGLE_CLOUD_PROJECT is unset ? the real Cloud Run dispatch needs a GCP
project id "
122                    "to resolve the job path; set it on the control plane (see
get_compute_backend)."
123                )
124            try:
125                from google.cloud import run_v2 # type: ignore[attr-defined]
126            except Exception as exc: # pragma: no cover - real SDK not present in CI
127                raise RuntimeError(
128                    "google-cloud-run is required for CloudRunCompute's real client; install
it on the "
129                    "control plane (it is intentionally NOT a CI/test dependency)."
130                ) from exc
131            client = (
132                run_v2.JobsClient()
133            ) # pragma: no cover - exercised only on the real M7 run
134            name = client.job_path(project, region, job_name)

```

```

135         overrides = run_v2.RunJobRequest.Overrides(
136             container_overrides=[
137                 run_v2.RunJobRequest.Overrides.ContainerOverride(args=argv)
138             ]
139         )
140         op = client.run_job(
141             request=run_v2.RunJobRequest(name=name, overrides=overrides)
142         )
143         return (
144             getattr(op, "metadata", None)
145             and getattr(op.metadata, "name", "")
146             or "execution"
147         )
148
149     def cancel_execution(self, execution: str) -> None:
150         # run_job falls back to the literal string "execution" when the operation
151         # metadata carries
152         # no name ? that is NOT a resolvable resource name. Fail loudly BEFORE the SDK
153         import
154         # (mirrors run_job's project guard, #342) so the misuse is exercisable in CI
155         (where
156         # google-cloud-run is absent).
157         if "/executions/" not in (execution or ""):
158             raise RuntimeError(
159                 f"not a fully-qualified Cloud Run execution name: {execution!r} ?
160                 expected "
161                 "projects/<p>/locations/<l>/jobs/<j>/executions/<e> (run_job could not
162                 resolve "
163                 "one from the operation metadata, so there is nothing to cancel by
164                 name)."
165             )
166         try:
167             from google.cloud import run_v2 # type: ignore[attr-defined]
168         except Exception as exc: # pragma: no cover - real SDK not present in CI
169             raise RuntimeError(
170                 "google-cloud-run is required to cancel a Cloud Run execution; install
171                 it on the "
172                 "control plane (it is intentionally NOT a CI/test dependency)."
173             ) from exc
174         client = (
175             run_v2.ExecutionsClient()
176         ) # pragma: no cover - exercised only on a real run
177         client.cancel_execution(request=run_v2.CancelExecutionRequest(name=execution))
178
179     def execution_terminal_state(self, execution: str) -> Optional[str]:
180         # Advisory status probe (see the ABC docstring): fetch the execution and
181         # classify it via the
182         # pure, unit-tested ``_classify_execution_state``. A non-resolvable name /
183         # absent SDK returns
184         # None (can't probe ? keep marker-polling) ? unlike run_job/cancel this never
185         # raises on bad
186         # input, because the probe is best-effort and must not fail a healthy run.
187         if "/executions/" not in (execution or ""):
188             return None
189         try:
190             from google.cloud import run_v2 # type: ignore[attr-defined]
191         except Exception: # pragma: no cover - real SDK not present in CI
192             return None
193         client = (
194             run_v2.ExecutionsClient()
195         ) # pragma: no cover - exercised only on a real run
196         exe = client.get_execution( # pragma: no cover - real run only
197             request=run_v2.GetExecutionRequest(name=execution)
198         )
199         return _classify_execution_state(exe) # pragma: no cover - real run only

```

```

190
191
192     def _classify_execution_state(exe: Any) -> Optional[str]:
193         """Map a Cloud Run ``run_v2.Execution`` to a terminal state label (``"Failed"`` /
194         ``"Cancelled"``
195         / ``"Succeeded"``), or ``None`` if it is still pending/running. PURE (no
196         SDK/network) so it is
197         unit-tested against both real proto-plus ``Execution`` objects and lightweight fakes
198         ? the actual
199         ``get_execution`` fetch is the only part that can't run in CI.
200
201         ``google-cloud-run`` is a **proto-plus** library, so a ``google.protobuf.Timestamp``
202         field is
203         marshalled to a native Python ``datetime`` when set and to ``None`` when unset ? NOT
204         to a proto
205         ``Timestamp`` with a ``.seconds`` field (a ``datetime`` has ``.second``, singular).
206         Terminality is
207         therefore "``completion_time`` is not None"; the per-task counts (plain ints) then
208         classify it,
209         with a cancel taking precedence (a cancelled execution reports ``cancelled_count >
210         0``)."""
211         if getattr(exe, "completion_time", None) is None:
212             return None # no terminal completion timestamp yet ? still pending/running
213         if getattr(exe, "cancelled_count", 0):
214             return "Cancelled"
215         if getattr(exe, "failed_count", 0):
216             return "Failed"
217         if getattr(exe, "succeeded_count", 0):
218             return "Succeeded"
219         return "Failed" # terminal but no positive success ? treat as failed (conservative)
220
221
222 class CloudRunCompute(ComputeBackend):
223     """Dispatch to a Cloud Run GPU Job + bucket-poll for completion."""
224
225     def __init__(
226         self,
227         *,
228         run_client: CloudRunJobClient,
229         job_name: str,
230         region: str,
231         project: Optional[str] = None,
232         storage_config: Optional[dict] = None,
233         policy: ExecutionPolicy = DEFAULT_POLICY,
234         token: Optional[str] = None,
235         container_overrides: Optional[tuple[str, ...]] = None,
236     ) -> None:
237         self._client = run_client
238         self._job_name = job_name
239         self._region = region
240         self._project = project
241         self._storage_config = storage_config or {}
242         self._policy = policy
243         # The dispatch token keys the done-marker so concurrent jobs to one bucket don't
244         collide.
245         # None ? a fresh uuid4 per run_infer call (prod); injected ? deterministic
246         (tests).
247         self._token = token
248         self._container_overrides = (
249             _DEFAULT_CONTAINER_OVERRIDES
250             if container_overrides is None
251             else tuple(container_overrides)
252         )
253
254     def run_infer(

```

```

245         self,
246         spec: ComputeJobSpec,
247         *,
248         on_wait: "Callable[[float], None] | None" = None,
249     ) -> ComputeResult:
250         out_root = spec.out_uri.rstrip("/")
251         token = self._token or uuid.uuid4().hex
252         done_uri = f"{out_root}/_dispatch/{token}.done"
253         argv: List[str] = [
254             "infer",
255             "--video-uri",
256             spec.video_uri,
257             "--out-uri",
258             spec.out_uri,
259             "--done-uri",
260             done_uri,
261         ]
262         if spec.segmenter:
263             argv += ["--segmenter", spec.segmenter]
264         if spec.skip_validation:
265             # The control-plane already validated + only dispatches on PASS ? tell the
worker not to
266             # re-decode and re-validate (redundant, and its config may resolve different
thresholds).
267             argv.append("--skip-validation")
268         for kv in (*spec.config_overrides, *self._container_overrides):
269             argv += ["--set", kv]
270
271         execution = self._client.run_job(
272             self._job_name, argv, region=self._region, project=self._project
273         )
274         LOG.info(
275             "cloud-run: dispatched job=%s exec=%s; bucket-polling %s",
276             self._job_name,
277             execution,
278             done_uri,
279         )
280
281         try:
282             marker = poll_until(
283                 lambda: self._poll_check(done_uri, str(execution or "")),
284                 self._policy,
285                 label="cloud-run-infer",
286                 logger=LOG,
287                 # Live heartbeat (#482): each still-waiting tick reaches the caller so
288                 # job.progress moves during the whole GPU run instead of freezing.
289                 on_tick=on_wait,
290             )
291         except PolicyTimeout as timeout_exc:
292             # RemoteExecutionFailed (the fast negative path in _poll_check) is NOT a
PolicyTimeout,
293             # so it propagates past here uncaught ? the dispatch handler maps it to a
distinct,
294             # clearly-labelled failure instead of the ~65-min opaque timeout this branch
handles.
295             marker = self._handle_poll_timeout(done_uri, str(execution), timeout_exc)
296             video_id = str(marker.get("video_id", ""))
297             # A MISSING returncode defaults to FAILURE (1), never success ? a present-but-
unreadable /
298             # garbled marker must not masquerade as rc=0. (The worker always writes both
fields.)
299             rc = int(marker.get("returncode", 1))
300             if not video_id:
301                 LOG.warning(
302                     "cloud-run: job=%s completion marker present but empty/unreadable (no "

```

```

303         "video_id) ? treating as failure",
304         self._job_name,
305     )
306     rc = rc or 1
307     outputs_uri = f"{out_root}/outputs/{video_id}/" if video_id else ""
308     report = self._read_json(f"{outputs_uri}report.json") if outputs_uri else {}
309     LOG.info(
310         "cloud-run: job=%s done (video_id=%s rc=%d)", self._job_name, video_id, rc
311     )
312     return ComputeResult(
313         returncode=rc,
314         outputs_uri=outputs_uri,
315         video_id=video_id,
316         report=report,
317         execution=str(execution or ""),
318     )
319
320 def _poll_check(self, done_uri: str, execution: str) -> Optional[dict]:
321     """One bucket-poll iteration, hardened with an Executions-API fast negative
322     path.
323
324     The done-marker is the SOLE success signal and is checked FIRST ? a present
325     marker always
326     wins (even if the execution's status later flips), so a worker that wrote a
327     SUCCESS *or a
328     FAILURE* marker resolves through the normal path with its real returncode. Only
329     when the
330     marker is absent do we consult the execution status: a worker that TERMINATED
331     without ever
332     writing a marker (crash / OOM / an old image argparse-aborting on an unknown
333     flag) would
334     otherwise make this poll spin the full ``remote_poll_max_wait_sec`` (~65 min)
335     before failing
336     actionable
337     opaquely. Detecting that terminal state collapses the whole class to a prompt,
338     failure.
339
340     Returns the marker dict (resolve the poll), ``None`` (keep waiting), or raises
341     ``RemoteExecutionFailed`` (abort the poll immediately ? the execution is over,
342     no marker
343     will ever appear)."""
344     marker = self._read_marker(done_uri)
345     if marker is not None:
346         return marker # success signal ? authoritative even if status later flips
347     state = self._terminal_execution_state(execution)
348     if state is None:
349         return None # still pending/running, or status unknown ? keep bucket-
350     polling
351     # The execution is terminal but no marker is present. Re-check the marker ONCE
352     # race where the worker wrote it between our two reads (mirrors the post-timeout
353     # _handle_poll_timeout). If it is genuinely absent, the process is gone and no
354     # EVER appear ? fail fast + diagnostically rather than wait out the whole poll
355     # budget.
356     marker = self._read_marker(done_uri)
357     if marker is not None:
358         return marker
359     raise RemoteExecutionFailed(
360         f"cloud-run: job={self._job_name} execution={execution or '<unknown>'}
361         reached "
362         f"terminal state {state!r} with NO completion marker ? the worker exited
363         before "
364         "writing one (crash / OOM / pre-marker abort, or a failed marker upload on

```

```

an "
352         "otherwise-Succeeded run). The execution is over, so the marker will never
appear; "
353         "inspect the Cloud Run Job execution log to diagnose: "
354         f"`gcloud run jobs executions describe {execution or '<execution>'} "
355         f"--region {self._region}`."
356     )
357
358     def _terminal_execution_state(self, execution: str) -> Optional[str]:
359         """Best-effort wrapper over the client's ``execution_terminal_state`` probe. A
raising or
360         flaky status API must NEVER kill a healthy wait (the done-marker stays the sole
authority),
361         so any exception ? and an empty/unresolvable execution name ? degrades to
``None`` (unknown
362         ? keep polling) and is DEBUG-logged, exactly today's marker-only behaviour."""
363         if not execution:
364             return None
365         try:
366             return self._client.execution_terminal_state(execution)
367         except Exception: # noqa: BLE001 ? advisory probe: never let it kill the
bucket-poll
368             LOG.debug(
369                 "cloud-run: execution status probe failed for %s ? treating as unknown",
370                 execution,
371                 exc_info=True,
372             )
373             return None
374
375     def _handle_poll_timeout(
376         self, done_uri: str, execution: str, timeout_exc: PolicyTimeout
377     ) -> dict:
378         """Poll deadline hit ? rescue a just-finished result, else cancel + re-raise.
379
380         The poll budget (``max_wait_sec``, default 1800s) can undercut the deployed
job's own
381         ``--task-timeout`` (3600s), so a slow-but-healthy worker may finish just past
the deadline:
382         ONE final marker check rescues that completed GPU run instead of wasting it. If
the marker
383         is genuinely absent, the execution is still running ? best-effort cancel it
(otherwise the
384         L4 keeps billing until ``--task-timeout``) and re-raise a ``PolicyTimeout``
carrying the
385         execution name so the operator can verify the state in the GCP console. A cancel
failure
386         NEVER masks the original timeout."""
387         try:
388             late = self._read_marker(done_uri)
389         except Exception: # noqa: BLE001 ? the final check is best-effort too: absence
? failure
390             LOG.warning(
391                 "cloud-run: final post-timeout marker check failed ? treating as
absent",
392                 exc_info=True,
393             )
394             late = None
395         if late is not None:
396             LOG.warning(
397                 "cloud-run: job=%s exec=%s done-marker found on the final post-timeout
check "
398                 "? rescuing the completed result",
399                 self._job_name,
400                 execution,
401             )

```

```

402         return late
403     try:
404         self._client.cancel_execution(execution)
405     except Exception as cancel_exc: # noqa: BLE001 ? best-effort: never mask the
timeout
406         cancel_note = f"best-effort cancel FAILED ({cancel_exc})"
407         LOG.warning(
408             "cloud-run: best-effort cancel of execution=%s failed: %s",
409             execution,
410             cancel_exc,
411             exc_info=True,
412         )
413     else:
414         cancel_note = "cancel requested"
415         LOG.info("cloud-run: requested cancel of execution=%s", execution)
416         raise PolicyTimeout(
417             f"{timeout_exc} [cloud-run job={self._job_name} "
418             f"execution={execution or '<unknown>'}; {cancel_note} ? verify the execution
"
419             "state in the GCP console]"
420         ) from timeout_exc
421
422     # --- storage-backed completion signal (bucket-poll) ---
423     def _read_marker(self, done_uri: str) -> Optional[dict]:
424         """The done-marker dict if the worker has written it, else None ? poll_until
keeps waiting."""
425         ref = parse_uri(done_uri)
426         if not get_storage(ref.backend, config=self._storage_config).exists(ref):
427             return None
428         return (
429             self._read_json(done_uri) or {}
430         ) # empty dict still terminates the poll (present == done)
431
432     def _read_json(self, uri: str) -> dict:
433         # Download to a real local dst FIRST: the cloud backends (gcs/s3) REQUIRE a dst
? a None dst
434         # raises ValueError ? unlike local's identity passthrough. Without an explicit
dst this read
435         # would silently return {} for every gs://|s3:// object (the only backends this
shim targets).
436         try:
437             ref = parse_uri(uri)
438             with tempfile.TemporaryDirectory(prefix="cr-read-") as td:
439                 dst = os.path.join(td, ref.name or "obj.json")
440                 local = fetch(uri, dst, config=self._storage_config)
441                 with open(local, "r", encoding="utf-8") as f:
442                     data: Any = json.load(f)
443                 return data if isinstance(data, dict) else {}
444         except Exception:
445             LOG.warning("cloud-run: could not read %s", uri, exc_info=True)
446             return {}
447

```

16. user (2026-07-08T05:06:00.017Z)

```

1     """Pluggable compute layer: in-process (default) + a Cloud Run GPU Job dispatch shim.
2
3     COMPUTE_DECOUPLED_SERVING M6. ``get_compute_backend(config)`` is the seam: it returns a
4     ``ComputeBackend`` ONLY when a non-local ``deployment.compute_target`` is selected, and
``None``
5     otherwise ? ``None`` means "run the worker's existing in-process path" (today's
behaviour, byte
6     identical). So the seam is **default-OFF**: with no profile / ``local_gpu`` /
``cpu_mock`` nothing
7     changes. Mirrors ``backend/storage``'s lazy factory idiom (cloud SDKs imported only when

```

```

needed).
8     """
9
10    from __future__ import annotations
11
12    import dataclasses
13    import os
14    from typing import Any, Optional
15
16    from backend.compute.base import (
17        ComputeBackend,
18        ComputeJobSpec,
19        ComputeResult,
20        RemoteExecutionFailed,
21    )
22    from backend.infrastructure.execution_policy import DEFAULT_POLICY, ExecutionPolicy
23
24    __all__ = [
25        "ComputeBackend",
26        "ComputeJobSpec",
27        "ComputeResult",
28        "RemoteExecutionFailed",
29        "get_compute_backend",
30    ]
31
32
33    def _storage_dict(config: Any) -> dict:
34        storage = getattr(config, "storage", None)
35        if storage is None:
36            return {}
37        if hasattr(storage, "model_dump"):
38            return storage.model_dump() # type: ignore[no-any-return]
39        return dict(storage) if isinstance(storage, dict) else {}
40
41
42    def get_compute_backend(
43        config: Any,
44        *,
45        run_client: Optional[Any] = None,
46        policy: ExecutionPolicy = DEFAULT_POLICY,
47        token: Optional[str] = None,
48        target_override: Optional[str] = None,
49    ) -> Optional[ComputeBackend]:
50        """Return the compute backend for ``config.deployment.compute_target``, or ``None``
to run
51        in-process (the default-OFF path).
52
53        Only ``compute_target == "cloud_run"`` returns a non-None backend (a
``CloudRunCompute``). Its
54        Cloud Run coordinates come from the control-plane env (``CLOUD_RUN_JOB`` /
``CLOUD_RUN_REGION``
55        / ``GOOGLE_CLOUD_PROJECT``) so no secrets live in config. ``run_client`` is
injectable (tests
56        pass a fake; production defaults to the real ``google-cloud-run`` client, owner-
verified on M7).
57
58        ``target_override`` lets a single request force a target (the SPA compute-picker
maps a
59        per-request ``compute="cloud"`` to ``target_override="cloud_run"``) without mutating
the
60        server's configured default. ``None`` falls back to
``config.deployment.compute_target``.
61        """
62        target = (
63            target_override

```



```

64         if target_override is not None
65         else getattr(getattr(config, "deployment", None), "compute_target", "") or ""
66     )
67     if target != "cloud_run":
68         return None # local_gpu / cpu_mock / "" ? run inline (today's path, unchanged)
69
70     # Lazy import so the cloud SDK shim is only pulled when a cloud target is actually
selected.
71     from backend.compute.cloud_run import CloudRunCompute, GoogleCloudRunJobClient
72
73     # Derive the poll budget from config UNLESS the caller injected an explicit policy ?
74     # deployment.remote_poll_max_wait_sec must outlast the worker Job's --task-timeout
so a
75     # slow-but-healthy paid GPU run is never abandoned mid-flight (CP-rebuild §4.3; the
old
76     # 1800s default undercut real 1080p60 jobs). Identity check on DEFAULT_POLICY keeps
an
77     # explicitly-passed policy (tests, callers with bespoke budgets) untouched.
78     if policy is DEFAULT_POLICY:
79         max_wait = getattr(
80             getattr(config, "deployment", None), "remote_poll_max_wait_sec", None
81         )
82         if max_wait:
83             policy = dataclasses.replace(DEFAULT_POLICY, max_wait_sec=float(max_wait))
84
85     job_name = os.environ.get("CLOUD_RUN_JOB", "rally-worker")
86     region = os.environ.get("CLOUD_RUN_REGION", "asia-south1")
87     project = os.environ.get("GOOGLE_CLOUD_PROJECT")
88     client = run_client if run_client is not None else GoogleCloudRunJobClient()
89     return CloudRunCompute(
90         run_client=client,
91         job_name=job_name,
92         region=region,
93         project=project,
94         storage_config=_storage_dict(config),
95         policy=policy,
96         token=token,
97     )
98

```

17. assistant (2026-07-08T05:06:05.692Z)

Let me read the compute base, config models, and find the compile job dispatch/API sites.

18. user (2026-07-08T05:06:07.211Z)

```

1     """The ``ComputeBackend`` ABC ? one interface for *where* the rally core runs.
2
3     COMPUTE_DECOUPLED_SERVING M6. Mirrors the ``StorageBackend`` ABC+factory idiom
4     (``backend/storage/``): a ``ComputeBackend`` decides **where** the pure
DETECT?WINDOW?STITCH
5     core executes ? in-process on the local box (truly-local) or on a **Cloud Run GPU Job**
6     (cloud-serving) ? without the core ever branching on the profile (D1).
7
8     The seam is **default-OFF**: only ``deployment.compute_target == "cloud_run"`` selects a
9     non-local backend. Every other target (``"" / ``local_gpu`` / ``cpu_mock``) keeps the
worker's
10    existing in-process path **byte-identical** ? the factory returns ``None`` (= "run
inline").
11    """
12
13    from __future__ import annotations
14
15    from abc import ABC, abstractmethod
16    from dataclasses import dataclass, field

```

```

17     from typing import Callable, Optional
18
19
20     class RemoteExecutionFailed(RuntimeError):
21         """A dispatched remote job's execution reached a TERMINAL state without ever writing
its
22         completion marker ? the worker crashed / OOM'd / aborted (e.g. an old image
argparse-rejecting
23         an unknown flag) before the marker push, or the marker upload itself failed on an
otherwise
24         Succeeded run.
25
26         Distinct from ``PolicyTimeout`` (poll budget exhausted while the execution may STILL
be
27         running): this says the execution is DEFINITELY over and no marker will ever
appear, so the
28         dispatcher must stop bucket-polling immediately instead of burning the full
29         ``remote_poll_max_wait_sec`` (~65 min) on a guaranteed-futile wait. The message
carries the
30         job + execution name + terminal state + the ``gcloud`` command to read the execution
log, so
31         the whole "worker crashed before the done-marker" class fails fast AND
diagnostically."""
32
33
34     @dataclass(frozen=True)
35     class ComputeJobSpec:
36         """One inference job, profile-agnostic. ``video_id`` is the MD5 the worker will
(re)compute
37         from the SAME input bytes, so the output always lands at
``<out_uri>/outputs/<video_id>/`` ?
38         deterministic regardless of where the job ran."""
39
40         video_uri: str # input (gs://? for cloud; a path/local:// for local)
41         out_uri: str # output ROOT (gs://? bucket for cloud)
42         segmenter: Optional[str] = None
43         # Raw worker ``--set k=v`` overrides. A tuple (not list/dict) so the spec stays
frozen/hashable.
44         config_overrides: tuple[str, ...] = ()
45         # Skip the worker's redundant re-validation. The DISPATCHER (the cloud-serving
control-plane) is
46         # the authoritative validation gate: it runs the SAME validators with ITS resolved
config and only
47         # dispatches on PASS, so a worker re-run is (a) redundant, (b) a second full-video
decode, and
48         # (c) outright WRONG when the worker's config resolves different thresholds ? it can
REJECT a clip
49         # the gate already passed. Default False keeps the direct ``worker infer`` CLI (no
control-plane)
50         # validating exactly as before.
51         skip_validation: bool = False
52
53
54     @dataclass(frozen=True)
55     class ComputeResult:
56         """The terminal state of a dispatched job. ``returncode`` mirrors the worker's exit
code
57         (0 ok / 2 validation-or-config / 3 segmenter-failure), plus **4 = remote dispatch
timeout or an
58         empty/unreadable completion marker** (the cloud path only). ``video_id`` is known
once the job's
59         done-marker is read; ``outputs_uri`` is where ``intervals.csv`` / ``report.json``
landed."""
60
61         returncode: int

```

```

62     outputs_uri: str = ""
63     video_id: str = ""
64     report: dict = field(default_factory=dict)
65     # The backend's execution identity for PROVENANCE (e.g. the Cloud Run execution
name) ? surfaced
66     # on the job result so "the cloud path actually ran" is independently verifiable
(gcloud run jobs
67     # executions describe <execution>). Empty for the in-process path / backends without
one.
68     execution: str = ""
69
70
71     class ComputeBackend(ABC):
72         """Backend-agnostic compute. ``run_infer`` dispatches ONE inference job and returns
only once
73         its outputs are durably available (bucket-poll for cloud), so a caller treats local
and cloud
74         execution identically."""
75
76         @abstractmethod
77         def run_infer(
78             self,
79             spec: ComputeJobSpec,
80             *,
81             on_wait: "Callable[[float], None] | None" = None,
82         ) -> ComputeResult:
83             """Run ``spec`` to completion and return its ``ComputeResult``.
84
85             ``on_wait(elapsed_sec)`` (optional) is invoked periodically while the backend is
86             still WAITING on the remote work ? the caller's live-progress heartbeat (#482),
so a
87             long GPU run doesn't look frozen. Implementations without a wait loop may ignore
it."""
88

```

19. user (2026-07-08T05:06:08.012Z)

```

1     """Pydantic models for the application configuration.
2
3     This module defines the canonical, validated, typed configuration for the
4     Badminton Highlight Indexer. Built-in defaults live here (single source of truth).
5
6     Unknown / typo'd config keys are rejected at load time via extra="forbid".
7     """
8
9     from __future__ import annotations
10
11     from typing import Any, List, Literal, Optional
12
13     from pydantic import BaseModel, Field, field_validator, model_validator
14
15     Sport = Literal["badminton", "tennis", "pickleball", "table_tennis", "cricket"]
16
17
18     class ValidationRelevanceConfig(BaseModel):
19         court_green_lower: list[int] = Field(default=[35, 40, 40])
20         court_green_upper: list[int] = Field(default=[85, 255, 255])
21         court_pixels_min_ratio: float = Field(default=0.05, ge=0.0, le=1.0)
22
23
24     class ValidationConfig(BaseModel):
25         min_resolution_width: int = 1280
26         min_resolution_height: int = 720
27         min_fps: float = 29.9
28         min_blur_threshold: float = 20.0

```

```

29     max_scene_cuts_per_minute: float = 0.5
30     pts_max_gap_seconds: float = Field(
31         default=10.0,
32         ge=0.0,
33         description=(
34             "Maximum allowed keyframe gap (seconds) before flagging as spliced. "
35             "GoPro H.264/H.265 recordings typically have GOP intervals of 1?4s; "
36             "a gap >10s suggests edited/missing content. Mobile phones, screen "
37             "recordings, and other encoders may use different GOP sizes ? tune "
38             "this threshold for your recording device."
39         ),
40     )
41     relevance: ValidationRelevanceConfig = Field(
42         default_factory=ValidationRelevanceConfig
43     )
44
45
46 class TrackNetConfig(BaseModel):
47     wsl_distro: str = "Ubuntu"
48     repo_dir: str = "~/models/TrackNetV4"
49     conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
50     conda_env: str = "TrackNetV4"
51     weights_path: str = ""
52     queue_length: int = 5
53     stage_in_wsl: bool = True
54     wsl_stage_dir: str = "~/clips"
55     timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
56     velocity_threshold: float = 0.02
57     # Cache hygiene: after a SUCCESSFUL run the staged video copy (and, for WASB, the
58     # decoded frame cache) are deleted automatically ? they are pure intermediates,
59     # regenerable from the source, and the dominant disk consumer (~GBs/video). Set
60     # true only for debugging. On FAILURE they are always kept so a re-run can resume.
61     keep_frames: bool = False
62     # Absolute FLOOR (GB) on WSL free space (at wsl_stage_dir) before a run, on top of
the
63     # ~15x-source headroom the frame-stage guard always enforces (D11 ? it now BLOCKS,
not warns).
64     # 0 = no extra floor (the headroom guard still applies).
65     min_free_gb: float = 0.0
66
67
68 class WasbIndexConfig(BaseModel):
69     """Typed config for the live WASB shuttle substrate (indexing.wasb).
70
71     Mirrors backend/pipeline/detectors/wasb_runner.py::WasbConfig.from_indexing_cfg so
72     these knobs actually take effect ? previously this whole block was silently dropped
73     because nested config sections defaulted to extra='ignore'.
74     """
75
76     wsl_distro: str = "Ubuntu"
77     repo_dir: str = "~/models/WASB-SBDT"
78     conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
79     conda_env: str = "wasb"
80     weights_path: str = (
81         "~/models/WASB-SBDT/pretrained_weights/wasb_badminton_best.pth.tar"
82     )
83     sport: str = "badminton"
84     wsl_stage_dir: str = "~/clips"
85     timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
86     velocity_threshold: float = 0.02
87     # --- Linux-native runner only (detector_impl='native'); ignored by the WSL runner.
---
88     # Env vars override these (WASB_DEVICE / WASB_PYTHON) for a config-free GPU box.
89     # torch device for the native runner. "auto" = cuda-if-available-else-cpu (the M4
CPU fallback

```

```

90      # for a no-GPU box); "cuda" stays strict (hard-fails without a GPU ? no silent slow-
CPU downgrade).
91      device: str = "cuda" # auto | cuda | mps | cpu
92      python_bin: str = (
93          "python" # the `wasb` conda env python (invoked by path, no `conda activate`)
94      )
95      # Cache hygiene: after a SUCCESSFUL run the staged video copy + decoded frame cache
96      # (the ~GBs/video disk hog) are deleted automatically ? pure intermediates,
regenerable
97      # from the source. Set true only for debugging. On FAILURE they are kept so a re-run
resumes.
98      keep_frames: bool = False
99      # Absolute FLOOR (GB) on WSL free space (at wsl_stage_dir) before a run, on top of
the
100     # ~15x-source headroom the frame-stage guard always enforces (D11 ? it now BLOCKS,
not warns).
101     # 0 = no extra floor (the headroom guard still applies).
102     min_free_gb: float = 0.0
103     # FAST PATH ? decode the video on the fly and feed frames straight to the detector,
104     # skipping the slow per-frame PNG extraction that dominates a long clip (~46k PNGs
for a
105     # 12-min 1080p60 video, which overran the 30-min timeout). Output is bit-identical
to the
106     # PNG path (PNG is lossless) ? VERIFIED on a real GPU 2026-06-20 (native stream ==
WSL disk,
107     # 1194/1195 frames byte-identical; native run-to-run byte-identical /
deterministic).
108     # Tri-state:
109     # None = per-runner default ? the WSL runner keeps DISK extraction
(unchanged Windows
110     # behaviour); the Linux-native runner (GPU box) defaults to STREAMING
(the perf win).
111     # True/False = force on/off for BOTH runners (e.g. set False for a container cv2
can't seek).
112     stream_video: Optional[bool] = None
113     # GPU-feed tuning for the native runner (measured on the first real L4 serving run,
114     # 2026-07-03: the default batch 8 with single-process loading left the GPU <1%
utilized ?
115     # the pipeline was CPU-feed-bound). None = wasb_infer.py's parity defaults (batch 8,
no
116     # prefetch ? byte-identical to the historical path). The cloud-serving preset opts
into
117     # the measured fast values; local/WSL behaviour is unchanged unless set explicitly.
118     batch_size: Optional[int] = None
119     # Depth of the bounded producer-thread queue that overlaps CPU-side batch assembly
120     # (decode + PIL + ToTensor) with GPU inference in streaming mode. None/0 = off.
121     prefetch_batches: Optional[int] = None
122     # GPU-RESIDENT DECODE (#506) ? DEFAULT OFF. "nvdec_resident" moves
decode+resize+normalize
123     # onto the GPU with no host round-trip: PyNvVideoCodec NVDEC decode ? affine
`grid_sample`
124     # resize replicating WASB's cv2.warpAffine EXACTLY (the #506 lesson: a naive
F.interpolate
125     # silently over-detects ? F1 0.58?0.07 ? and only shows on hard clips) ? GPU
ImageNet
126     # normalize ? HRNet. Validated on a real L4 (GX010128, SERVED preset, 2026-07-06):
~1.7x
127     # fps (35.8?60.2), GPU util 47?79% / CPU 62?35%, F1 0.574 ? baseline 0.582. Native
(Linux)
128     # runner only ? needs the torch-2.x + PyNvVideoCodec WASB env
(deploy/cloudrun/Dockerfile)
129     # with libnvcuvid/libnvidia-encode present, and device=cuda; the WSL runner ignores
it.
130     # Flip gate: the full 15-clip golden set at the SERVED preset (GX010128 must pass).
131     decode_backend: Literal["cpu", "nvdec_resident"] = "cpu"

```

```

132
133     @model_validator(mode="after")
134     def _nvdec_resident_requires_cuda(self) -> "WasbIndexConfig":
135         # NVDEC decode is CUDA-only; reject the contradictory combo at config load
instead of
136         # failing deep inside a paid GPU run. "auto" is allowed here ? the native
runner's
137         # healthcheck hard-fails if auto resolves to cpu with this backend on.
138         if self.decode_backend == "nvdec_resident" and self.device not in ("cuda",
"auto"):
139             raise ValueError(
140                 f"indexing.wasb.decode_backend='nvdec_resident' requires device 'cuda'
or "
141                 f"'auto' (got {self.device!r}) ? NVDEC decode runs on the GPU."
142             )
143         return self
144
145     @field_validator("batch_size")
146     @classmethod
147     def _batch_size_positive(cls, v: Optional[int]) -> Optional[int]:
148         # A DataLoader batch must be >=1; reject at config time so a typo can't reach
149         # `wasb_infer.py --batch-size` as an invalid value (the runner only passes the
flag
150         # when truthy, so 0 stays "off/parity-default").
151         if v is not None and v < 1:
152             raise ValueError(f"indexing.wasb.batch_size must be >= 1 when set (got
{v})")
153         return v
154
155     @field_validator("prefetch_batches")
156     @classmethod
157     def _prefetch_batches_non_negative(cls, v: Optional[int]) -> Optional[int]:
158         # 0 = off (the default). Reject negatives: the runner passes any truthy value as
159         # `--prefetch-batches`, and `_PrefetchIter` clamps depth with max(1,
int(depth)), so a
160         # negative would SILENTLY enable prefetch at depth 1 instead of failing loudly.
161         if v is not None and v < 0:
162             raise ValueError(
163                 f"indexing.wasb.prefetch_batches must be >= 0 when set (got {v})"
164             )
165         return v
166
167
168     class FusionConfig(BaseModel):
169         """Config for the multi-signal `fusion_hybrid` segmenter (MULTI_SIGNAL_FUSION_PLAN
P2).
170
171         Every default reproduces control behaviour: the fusion segmenter persists the
172         shuttle net-crossing count (Gate A signal) and ? only when ``cue_flags['person']``
is
173         true ? the person-cue features, but it does NOT gate the served handoff unless a
174         threshold is raised (``min_crossings``/``min_players`` default 0 = OFF). The court
mask
175         is optional: ``court_polygon=None`` ? Gate B counts every detection (player-count-
only);
176         supplied ? off-court persons (spectators / adjacent courts) are excluded.
177         """
178
179         # Which trajectory substrate seeds the windows (shuttle stays primary).
180         substrate: Literal["wasb", "tracknet"] = "wasb"
181         # Windowing velocity threshold for the fusion family (the engine reads
182         # indexing.<family>.velocity_threshold). Default matches the substrates' 0.02; set
183         # equal to indexing.wasb.velocity_threshold to A/B fusion against a tuned wasb run.
184         velocity_threshold: float = 0.02
185         # Per-cue enable flags, e.g. {"person": true}. Person cue is OFF unless explicitly

```

```

186         # enabled ? so the default path never imports torch / touches the GPU.
187         cue_flags: dict[str, bool] = Field(default_factory=dict)
188         # Served Gate A: drop windows whose shuttle net-crossings < this from the AI
handoff.
189         # 0 = OFF (no gating; persisted candidates are always the full superset for offline
re-scoring).
190         min_crossings: int = Field(default=0, ge=0)
191         # Served Gate B: when the person cue is on, drop windows with median in-court
players
192         # < this. 0 = OFF.
193         min_players: int = Field(default=0, ge=0)
194         # Optional court mask as normalised 0-1 [[x, y], ...] polygon. None = no mask.
195         court_polygon: Optional[list[list[float]]] = None
196         # Net line for the person two-sided-motion split (person features only ? the shuttle
197         # net-crossing gate keeps rally_gate's camera-agnostic auto-estimate).
198         net_axis: Literal["x", "y"] = "y"
199         net_position: float = Field(default=0.5, ge=0.0, le=1.0)
200
201
202     class IndexingConfig(BaseModel):
203         default_segmenter: str = "yolo_hybrid"
204         use_gpu: bool = True
205         # Detector backend for the trajectory hybrids (wasb/tracknet). "auto" = the real
206         # WSL runner (default, production); "stub" = a deterministic CPU mock (no GPU/WSL)
207         # so the whole pipeline is regression-testable on a CPU box / CI; "native" = the
208         # Linux-native WASB runner (no WSL/conda-activate; a GPU box ? M3.5). The
209         # RALLY_DETECTOR_IMPL env var overrides this. See
pipeline/detectors/{stub,native_wasb}_runner.py.
210         detector_impl: str = "auto"
211         motion_threshold: float = 0.08
212         min_segment_duration: float = 3.0
213         dead_time_merge_gap: float = 2.0
214         chunk_duration: float = 90.0
215         chunk_overlap: float = 10.0
216         detector_model: str = "fasterrcnn_resnet50_fpn"
217         detector_score_threshold: float = 0.5
218         yolo_velocity_threshold: float = 0.15
219         yolo_frame_skip: int = 3
220         yolo_tracker: str = "bytetrack.yaml"
221         yolo_prefetch_workers: int = 2
222         min_action_window_duration: float = 2.0
223         # BOUNDED WINDOWING (over-merge fix W1, RALLY_QUALITY_RESEARCH.md §6). 0 = OFF; > 0
= a window
224         # longer than this many seconds is recursively split at its largest internal
inactivity gap (?
225         # ``window_min_split_gap`` s ? the most likely rally boundary), so one window can't
swallow
226         # several rallies + the dead time between them. Never merges, only cuts (recall
can't drop).
227         # *** R1 MILESTONE DEFAULT-ON (cap=6): the smallest safe local-windowing flip.
Measured n=13
228         # golden, R2 tolF1: baseline?milestone (cap=6 + R4 below) ?+0.222, sign 12/0,
p=0.0005; cap=6
229         # beats cap=12 ?+0.110, sign 12/0; net over-seg guard PASS (mean merge_rate
0.651?0.161). ***
230         max_window_duration: float = 6.0
231         window_min_split_gap: float = (
232             0.5 # min internal gap (s) that qualifies as a split point
233         )
234         # HARD-CAP fallback for W1: when True, an over-long window with NO internal
inactivity gap
235         # (a continuously-active trajectory ? e.g. the real-footage mahadevpura-1 174s blob)
is
236         # chopped at its midpoint until ? max_window_duration, making the cap a true upper
bound.

```

```

237         # Only cuts (recall can't drop). OFF by default until measured/promoted.
238         window_hard_cap: bool = False
239         # R4 rally-END inactivity trim (s, 0 = OFF). After a rally the shuttle keeps moving
slowly
240         # (picked up / walked) above velocity_thresh, so the window over-extends its END
(the measured
241         # [|?end| gap vs golden). Trim the post-rally tail back to the last frame whose
trailing
242         # window stays dense with FAST motion (velocity > inactivity_rally_thresh). Only
cuts.
243         # *** R1 MILESTONE DEFAULT-ON (1.5/0.20/0.30): on top of cap=6, R4 adds a small but
significant
244         # end-precision gain ? ?tolF1 +0.040, sign 9/1/3, p=0.022, [|?end| 4.96?3.31s (n=13).
The W1 cap
245         # is the dominant lever; R4 is the marginal increment. See QUALITY_ITERATIONS.md
"R1". ***
246         window_inactivity_end: float = 1.5
247         inactivity_rally_thresh: float = 0.20
248         inactivity_min_density: float = 0.30
249         # R4 density DENOMINATOR fix (code-review finding B). The R4 trim's fast-fraction
denominator is
250         # the COUNT of active samples in the trailing window, which sparse tracking inflates
(one fast
251         # sample after a gap ? 100% dense ? tail held open). When True, the denominator is
the window's
252         # TEMPORAL length instead. No-op under dense tracking (sample-count == temporal-
span), so it only
253         # corrects the sparse-track case. *** P2b PROMOTED DEFAULT-ON (2026-06-29): golden-
set A/B at the
254         # SERVED operating point (n=15, 595 rallies) ? tolF1 0.353?0.364 (?tolF1 +0.010,
sign 10W/4L, NO
255         # regression), [|?end| 2.97?2.72s, seg_count_ratio ?0.103 (p=0.006), split_rate
?0.010 (p=0.016),
256         # R1 over-seg guard PASS; served_regresses(eps=0.0)=False ?
promote_if_served_safe(structural)=True.
257         # See docs/RALLY_DETECTION_FIXES_PLAN.md §7/§8. ***
258         inactivity_temporal_density: bool = True
259         # R5 blob-fallback (default OFF). LAST-RESORT splitter for the continuously-active
blob: after
260         # the W1 gap-split AND the R4 inactivity trim have each had their chance, a group
STILL longer
261         # than window_blob_factor × max_window_duration (no internal gap, no rally-end
signal ? e.g.
262         # mahadevpura-1's ~65s residual ? 11× a 6s cap) is midpoint-chopped to ? the cap as
a FORCED cut
263         # (logged + flagged on the window) so the degenerate single-window outcome is
avoided and the
264         # clip is surfaced as needing rally-STATE cues, not a bigger cap. Differs from
window_hard_cap
265         # (which chops BEFORE R4 and over-segments normal clips): blob-fallback runs AFTER
R4 and only
266         # force-cuts the genuine blob (the factor gate). Only cuts (recall can't drop).
267         window_blob_fallback: bool = False
268         # Magnitude gate that keeps R5 surgical: only a group longer than this ×
max_window_duration is a
269         # blob (a normal long rally ~1-2× the cap is left alone; the mahadevpura-1 residual
is ~11×).
270         # Default 4.0 is do-no-harm on the golden corpus (rescues only the continuously-
active
271         # clips, every other clip a no-op within noise) ? see docs/QUALITY_ITERATIONS.md
"R5".
272         window_blob_factor: float = 4.0
273         log_yolo_candidates: bool = True
274         store_yolo_candidates: bool = True
275         ai_handoff_padding: float = 2.0

```



```

276     ai_max_workers: int = 5
277     # Width (px) the gemini segmenter re-encodes chunks to before upload. Stream-copied
278     # chunks of a high-bitrate (4K) source blow past the provider upload cap
279     # (backend/providers/gemini.py::MAX_UPLOAD_MB) and are refused ? observed live as
280     # "0 segments / success". Re-encoding also makes chunk boundaries frame-accurate
281     # (stream copy cuts at the nearest keyframe). 0 = legacy stream-copy at source res.
282     # Matches gemini_refine's --clip-scale-width default.
283     ai_chunk_scale_width: int = 1280
284     # FULLY-LOCAL mode (default OFF): when true, the trajectory-hybrid segmenters
285     # (wasb_hybrid / tracknet_hybrid / fusion_hybrid) skip the Phase-3 AI handoff
entirely and
286     # emit their local candidate windows AS the rallies ? no provider, no API key,
nothing
287     # uploaded. Lets the local detector run standalone (e.g. to benchmark it against the
Gemini
288     # path, or to operate with zero cloud dependency). The pure `gemini` segmenter
ignores this.
289     skip_ai_handoff: bool = False
290     # Optional debug knob: trim every video to the first N seconds before processing.
291     debug_duration: Optional[float] = None
292
293     tracknet: TrackNetConfig = Field(default_factory=TrackNetConfig)
294     wasb: WasbIndexConfig = Field(default_factory=WasbIndexConfig)
295     fusion: FusionConfig = Field(default_factory=FusionConfig)
296
297     @field_validator("default_segmenter")
298     @classmethod
299     def segmenter_must_be_registered(cls, v: str) -> str:
300         # Lazy import: the segmenter registry imports modules that may import
301         # config, so we resolve it at validation time rather than at module load.
302         from backend.pipeline.segmenters.base import segmenter_registry
303
304         if segmenter_registry.get(v) is None:
305             available = sorted(segmenter_registry.list_available().keys())
306             raise ValueError(
307                 f"Segmenter '{v}' is not registered. Available: {' ', ' '.join(available)}"
308             )
309         return v
310
311     @model_validator(mode="after")
312     def _chunk_step_must_be_positive(self) -> "IndexingConfig":
313         # The chunked segmenters advance by (chunk_duration - chunk_overlap); a non-
positive
314         # step makes `while t < duration: t += step` loop forever, growing memory before
any
315         # GPU work. Reject the bad config at load time instead.
316         if self.chunk_overlap >= self.chunk_duration:
317             raise ValueError(
318                 f"indexing.chunk_overlap ({self.chunk_overlap}) must be < chunk_duration
"
319                 f"({self.chunk_duration}); otherwise chunking never advances."
320             )
321         return self
322
323
324     class OutputConfig(BaseModel):
325         default_highlights_name: str = "highlights.mp4"
326         db_name: str = "sports_indexer.db"
327         # Directory where the DB, highlight reels, and processed clips are written.
328         # The CLI's --output-dir overrides this at startup (see backend/main.py::main).
329         output_dir: str = "output"
330
331
332     class WatermarkConfig(BaseModel):
333         """Branding overlay burned into each highlight clip during the per-clip re-encode

```

```

334 (backend/pipeline/stitcher/compiler.py). **On by default** ? set `enabled: false`
335 (under `stitching.watermark` in config.json) to turn it off. The text is fully
336 configurable. The concat stage stays stream-copy (-c copy)."""
337
338 enabled: bool = True
339 text: str = "Generated by Khelsutra.guru"
340 logo_path: str = "" # optional transparent PNG overlay; "" = no logo
341 font_path: str = "" # optional .ttf; "" = use the fontconfig `font` family below
342 font: str = "Sans" # fontconfig family used when font_path is empty
343 font_size_ratio: float = Field(
344     default=0.035, gt=0.0, le=0.5
345 ) # text height as a fraction of frame height
346 opacity: float = Field(default=0.85, ge=0.0, le=1.0)
347 box: bool = True # faint backing box behind the text for legibility
348 margin: int = Field(default=24, ge=0) # px inset from the chosen corner
349 position: Literal["bottom-right", "bottom-left", "top-right", "top-left"] = (
350     "bottom-right"
351 )
352
353
354 class StitchingConfig(BaseModel):
355     begin_padding: float = 0.5
356     end_padding: float = 0.5
357     use_cache: bool = False
358     min_shot_count: int = 1
359     # Long-rally reel filter (seconds; 0 = OFF). Drop detected rally pairs whose
duration
360     # (end ? start) is below this from the highlight reel, BEFORE compiling. A cleaner
361     # reel-selection knob than `min_shot_count`: shot_count is unlabeled / "Unknown" on
the
362     # fully-local no-AI path, whereas duration is derived from the rally's own
start/end. The
363     # local detector over-fires and its false positives are mostly SHORT (motion blips,
364     # background-court fragments), so this keeps the reel to the eye-catching long
rallies.
365     # OFF by default ? the detector output is unchanged, only which rallies reach the
reel.
366     min_rally_duration: float = Field(default=0.0, ge=0.0)
367     watermark: WatermarkConfig = Field(default_factory=WatermarkConfig)
368
369
370 class LoggingConfig(BaseModel):
371     """Logging knobs for the run entrypoints (see backend/utils/logging_setup.py)."""
372
373     # Third-party HTTP/RPC client loggers (httpx, httpcore, urllib3, google-genai,
374     # google.auth, grpc) emit one INFO line per request ? dozens per Gemini run, which
375     # buries our own [rally]/[compile] lines in run.log during manual QA. Default raises
376     # them to WARNING; set true to restore full chatter for request-flow debugging.
377     verbose_http: bool = False
378
379
380 class SecurityConfig(BaseModel):
381     # When set, /api/process video_path must resolve under this directory (hosted/tenant
382     # mode). Default None preserves the single-user local 'any local file' workflow.
383     allowed_video_root: Optional[str] = None
384
385     # --- Chunked upload (B2, PR-2) ---
386     # Landing zone for browser uploads (assembled from parts). Per-user subdirectories
387     # arrive with PR-3 identity scoping. ? In hosted mode, set allowed_video_root to
388     # contain upload_dir or /api/process will reject the uploaded paths.
389     upload_dir: str = "output/uploads"
390     # Whole-file cap (default 4 GB) and per-part cap. Parts stay under ~100 MB because
391     # free-tier edge proxies cap request bodies there (HOSTING_PLAN §2).
392     max_upload_mb: int = 4096
393     max_part_mb: int = 95

```

```

394     allowed_upload_extensions: List[str] = ["mp4", "mov"]
395
396     # --- M9 edge auth (Cloudflare Access, D9) ---
397     # DEFAULT-OFF: edge_auth_enabled=False ? no JWT is required/verified (the local
single-user
398     # path is unchanged). When True, /api/process requires + verifies the `Cf-Access-
Jwt-Assertion`
399     # header against the team JWKS (so a direct hit to the Cloud Run URL can't spoof the
identity
400     # header) and tags the job with the verified user_id (owner_id). The team domain
(e.g.
401     # ``https://<team>.cloudflareaccess.com``) supplies the JWKS + issuer; the AUD is
the CF Access
402     # application's audience tag. NO secrets here ? these are public identifiers.
403     edge_auth_enabled: bool = False
404     cf_team_domain: str = (
405         "" # e.g. https://<team>.cloudflareaccess.com (issuer + /certs JWKS)
406     )
407     cf_access_aud: str = "" # the CF Access application AUD tag the token must carry
408     # (CORS allow-origins is set via the RALLY_CORS_ALLOW_ORIGINS env var ? read at
import in
409     # backend/main.py so importing the app does no filesystem I/O; default ["*"].
Credentials stay
410     # OFF ? Cf-Access is a header, not a cookie.)
411
412
413     class StorageConfig(BaseModel):
414         """Pluggable storage backend (docs/DECOUPLED_COMPUTE/). Default ``local`` = today's
415         behaviour (filesystem paths, byte-for-byte). ``s3`` covers AWS + Cloudflare R2 + DO
Spaces
416         + MinIO (via ``endpoint_url``); ``gcs`` = Google Cloud Storage; ``gdrive`` = a
locally MOUNTED
417         Google Drive (``gdrive.mount_root``). Per-backend settings are loose dicts passed to
the constructor ?
418         **secrets are never stored here**, only endpoints / regions / file paths / source-
selectors
419         (boto3/google auth chains supply credentials). See ``backend/storage/``. """
420
421         backend: Literal["local", "s3", "gcs", "gdrive", "gdrive-api"] = "local"
422         # Output root for produced artifacts. "" => fall back to output.output_dir (non-
breaking).
423         output_root: str = ""
424         # --- Input side (COMPUTE_DECOUPLED_SERVING M2; parallel to the output
`backend`/`output_root`).
425         # Default ``local`` / "" keeps today's behaviour: a bare path / ``local://`` URI is
read IN PLACE
426         # (identity, no copy). Set these to read inputs from a bucket / mounted Drive ? a
bare/relative
427         # input name is resolved against ``input_root`` (e.g.
``input_root``="gs://bucket/videos"`` + "x.mp4"
428         # => ``gs://bucket/videos/x.mp4``) or, if ``input_root`` is empty, prefixed with
``input_backend``
429         # (``"gcs"`` + "bucket/x.mp4" => ``gcs://bucket/x.mp4``). An explicit scheme on the
input
430         # (``gs://?``/``s3://?``) or an absolute local path ALWAYS wins. Non-local inputs
are fetched to a
431         # scratch dir before processing (see ``backend.storage.acquire_local_input``).
432         input_backend: Literal["local", "s3", "gcs", "gdrive", "gdrive-api"] = "local"
433         input_root: str = ""
434         # --- Per-video workspace (docs/PER_VIDEO_WORKSPACE.md) ---
435         # One folder per video keyed by the MD5 video_id:
<root>/[<workspace_prefix>]/<video_id>/...
436         # so a local deploy and a cloud (bucket) deploy share ONE relative layout (only the
root
437         # differs). DEFAULT-OFF: convergence is phased + flipped via a Launch Report (D5).

```

```

When OFF
438     # the legacy flat layout is used unchanged.
439     single_folder_per_video: bool = False
440     # Durable workspace root URI. "" => precedence: output_root, then output.output_dir.
441     workspace_root: str = ""
442     # Cloud namespace prefix under the root so per-video folders sit tidily beside
443     # videos/ labels/ weights/ (e.g. gs://bucket/workspaces/<video_id>/). "" = root
directly.
444     workspace_prefix: str = "workspaces"
445     local: dict[str, Any] = Field(default_factory=dict)
446     s3: dict[str, Any] = Field(default_factory=dict)
447     gcs: dict[str, Any] = Field(default_factory=dict)
448     gdrive: dict[str, Any] = Field(default_factory=dict)
449     # gdrive-api settings (e.g. {"root_folder_id": "<id>"}); read via the
hyphen?underscore lookup in
450     # backend.storage.get_storage. NO secrets ? creds come from ADC over the Drive
scope.
451     gdrive_api: dict[str, Any] = Field(default_factory=dict)
452
453
454     class DeploymentConfig(BaseModel):
455         """Deployment profile ? one switch that selects a coherent bundle of compute +
storage
456         knobs (COMPUTE_DECOUPLED_SERVING M1, ``docs/COMPUTE_DECOUPLED_SERVING/README.md``
$4.3).
457
458         **Pure / additive:** the empty default profile (``""``) applies NO preset, so the
pipeline
459         behaves exactly as before this section existed. A profile is opt-in (``config.json``
460         ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); auto-detect only ever
461         *suggests* one (D15 ? cloud spend is never entered silently). Nothing in the core
reads
462         these fields yet (M1): selecting a profile works purely by pre-setting the EXISTING
knobs
463         (``indexing.detector_impl`` / ``skip_ai_handoff``, ``storage.backend``) in the
loader.
464
465         Valid values mirror ``backend/config/presets.py`` (parity-pinned by a unit test)."""
466
467         # "" = no profile (today's manual knobs). Others apply presets.PROFILE_PRESETS.
468         profile: Literal["", "truly-local", "cloud-serving", "cpu-mock"] = ""
469         # Where the core runs. Set by the profile preset; overridable in config.json. Inert
in M1
470         # (recorded for observability; consumed by the ComputeBackend seam from M2+).
471         compute_target: Literal["", "local_gpu", "cloud_run", "cpu_mock"] = ""
472         # Cloud-serving throughput cap (#466): how many DETECT jobs the control-plane drain
may dispatch +
473         # bucket-poll CONCURRENTLY when compute_target=="cloud_run". Each is an I/O-bound
Cloud Run Job
474         # execution (the GPU work is on ephemeral L4s, WASB-only has no Gemini), so ? unlike
a truly-local
475         # single box ? they need not serialize. Bounds concurrent PAID L4 executions. Non-
cloud_run targets
476         # ignore this and stay serial (max_workers=1), preserving the truly-local GPU/Gemini
serialization.
477         max_concurrent_remote_jobs: int = Field(default=4, ge=1, le=64)
478         # Poll budget for one dispatched remote job (CONTROL_PLANE_DURABLE_REBUILD $4.3).
Must sit
479         # ABOVE the worker Job's own --task-timeout (3600s) so the worker's timeout stays
the
480         # authority and the dispatcher's one-shot post-timeout marker rescue still applies ?
the old
481         # 1800s ExecutionPolicy default undercut real 1080p60 jobs (~22 min detect +
queue/cold-start,
482         # observed on the 2026-07-05 owner CUJ) and could abandon a slow-but-healthy paid

```

```

GPU run.
483         remote_poll_max_wait_sec: float = Field(default=3900.0, gt=0)
484
485
486     class BillingConfig(BaseModel):
487         """Per-job cost ledger (COMPUTE_DECOUPLED_SERVING M8, D8).
488
489         **DEFAULT-OFF:** ``enabled=False`` ? nothing records cost (the v6 cost columns stay
NULL =
490         today's behaviour). When enabled, a job's cost is computed + stored in **USD** via
the versioned
491         ``pricebook`` DB table and **displayed in INR** at ``fx_inr_per_usd``. The default
pricebook
492         version is the seeded ``placeholder-v0`` (every rate ``0.0`` ? cost ``0``) until the
owner
493         inserts a real, calibrated version and points ``pricebook_version`` at it.
494
495         ? Even with ``enabled=True`` + a real pricebook, recorded cost stays **0** until a
follow-up
496         wires the worker to emit per-job ``usage`` (gpu/wall/egress/gemini) ? M8 lands the
ledger +
497         pricebook + the recording seam; usage emission (the metering hooks) is a separate
step."""
498
499         enabled: bool = False
500         pricebook_version: str = (
501             "placeholder-v0" # the seeded $0 version; owner flips to a real one
502         )
503         fx_inr_per_usd: float = 83.0 # USD?INR display rate (Bangalore market)
504         margin: float = 0.0 # resale markup (?0); 0 = bill at cost
505
506
507     class FlywheelConfig(BaseModel):
508         """Data-flywheel harness (COMPUTE_DECOUPLED_SERVING M11, D12) ? the
consent?capture?eval?goldenize
509         loop: on a CONSENTED adult job, durably **capture** the source + rally output into
the per-video
510         workspace ? **shadow-eval** owned-vs-Gemini ? **propose promotion** to the golden
TRAINING set, so
511         the owned model climbs toward the D11 ship floor.
512
513         **DEFAULT-OFF + INERT:** ``enabled=False`` ? nothing is captured/evaluated (byte-
identical to
514         today). ? **Consent is FAKED-DENY until counsel + the consent schema land:** even
``enabled=True``
515         captures **nothing**, because ``flywheel.consent_attested`` currently returns
``False`` for every
516         job (the privacy guardrail ? D10: adults-only, explicit opt-in, derived-data + auto-
delete-raw).
517         This is the PLUMBING only; the real consent gate (counsel-reviewed ToS/DPA ? consent
columns ?
518         wiring) + the owned-model + golden-label data needed to *activate* shadow-
eval/promotion are
519         owner/counsel/GPU scope. Pairs with M9-consent."""
520
521         enabled: bool = False
522         workspace_prefix: str = (
523             "flywheel" # per-video workspace prefix for captured artifacts
524         )
525         shadow_eval: bool = True # run the owned-vs-Gemini shadow comparison on capture
526         min_promote_n: int = 5 # promote_if_served_safe `min_n` (the served-safe gate)
527
528
529     class AppConfig(BaseModel):
530         """Root configuration object.

```

```

531
532     All runtime configuration should be accessed via an instance of this class
533     (or the loader) rather than raw dict lookups.
534     """
535
536     sport: Sport = "badminton"
537     ai_provider: str = "gemini"
538     # Auto-tracking alias: 2.5-flash is legacy now that 3.5-flash is the stable
flagship;
539     # the alias never deprecates underneath us. Eval reports record the run date
alongside.
540     ai_model: str = "gemini-flash-latest"
541
542     validation: ValidationConfig = Field(default_factory=ValidationConfig)
543     indexing: IndexingConfig = Field(default_factory=IndexingConfig)
544     output: OutputConfig = Field(default_factory=OutputConfig)
545     stitching: StitchingConfig = Field(default_factory=StitchingConfig)
546     security: SecurityConfig = Field(default_factory=SecurityConfig)
547     logging: LoggingConfig = Field(default_factory=LoggingConfig)
548     storage: StorageConfig = Field(default_factory=StorageConfig)
549     deployment: DeploymentConfig = Field(default_factory=DeploymentConfig)
550     billing: BillingConfig = Field(default_factory=BillingConfig)
551     flywheel: FlywheelConfig = Field(default_factory=FlywheelConfig)
552
553     # Reject unknown keys ? typed config migration is complete (D3).
554     model_config = {
555         "extra": "forbid",
556         "validate_assignment": True,
557     }
558
559
560     # Convenience type alias for callers that want to annotate
561     Config = AppConfig
562

```

20. user (2026-07-08T05:06:09.125Z)

```

36:# work: CPU/GPU-heavy in-process detection + the compile stitch serialize via
38:# and compile/local-override jobs stay serial. The count is fixed at first _get_executor()
creation.
75:    OVERLAPS. This does NOT parallelize CPU/GPU-heavy LOCAL work ? in-process detection and
the compile
77:    single-box GPU + rate-fragile Gemini invariant holds, and compile/local-override jobs
stay serial.
157:    # Dispatch by kind: 'compile' (#329, the async reel stitch) vs 'process' (NULL/default ?
the
159:    # jobs feed the M8 cost ledger + M11 flywheel (a compile has no detection
usage/segments).
160:    is_compile = job.get("kind") == "compile"
162:        if is_compile:
163:            result = _main._execute_compile(
180:            if not is_compile:

```

21. assistant (2026-07-08T05:06:15.520Z)

Now let me read the job worker dispatch, the compile API endpoint, and the VideoCompiler/ffmpeg stitcher.

22. user (2026-07-08T05:06:16.922Z)

```

1     """Async job worker ? drain/wake loop (DEFERRED_HARDENING_PLAN #16, EXECUTION_POLICY.md
P3).
2
3     Extracted verbatim from backend.main to shrink the orchestrator; behavior is unchanged.

```

```

4     The worker DRAINS the jobs table via `claim_due_job` and honours WAIT_AND_RESUME: a job
5     that raises `Jobwaiting` is parked (status 'waiting' + `next_poll_at` + progress) and
6     re-claimed when due ? with no central scheduler (the DB is the clock).
7
8     The monkeypatch contract these functions rely on (and that tests depend on):
9     every cross-function seam ? `db_instance`, `global_config`, `_execute_processing`,
10    `_arm_wake_timer`, `_get_executor`, `_MAX_DRAIN_WAKE_SEC` ? is resolved THROUGH the
11    `backend.main` module object at call-time (`import backend.main as _main`), never as a
12    bare local. `backend.main` re-exports these names, so `monkeypatch.setattr(backend.main,
13    ...)`` still intercepts the inter-function calls (e.g. a parked job arms the `*patched*`
14    `_arm_wake_timer`, so no real daemon timer leaks in tests). Importing `backend.main`
    lazily
15    inside the functions also avoids the circular import at module load (main.py imports
    this
16    module; this module must not import main at import time).
17    """
18
19    import json
20    import logging
21    import os
22    import threading
23    import traceback
24    from concurrent.futures import ThreadPoolExecutor
25    from typing import Any, Dict, Optional
26
27    logger = logging.getLogger(__name__)
28
29    # Async job executor (DEFERRED_HARDENING_PLAN #16). max_workers=1 by default: a single
30    # self-hosted box serializes GPU work, and the Gemini provider is rate-fragile under
31    # parallel uploads (2026-06-10 ops learning: serial uploads + backoff). The segmenters
32    # already parallelize their AI handoff internally via their own pools.
33    # EXCEPTION (#466): when the server CAN dispatch to Cloud Run (see
    _cloud_dispatch_possible), a REMOTE
34    # job is only a bucket-poll (I/O-bound; GPU is on ephemeral L4s, WASB-only has no
    Gemini), so the drain
35    # scales to deployment.max_concurrent_remote_jobs to overlap those. This does NOT
    parallelize local
36    # work: CPU/GPU-heavy in-process detection + the compile stitch serialize via
37    # orchestration._LOCAL_COMPUTE_LANE regardless of the worker count, so the single-box
    invariant holds
38    # and compile/local-override jobs stay serial. The count is fixed at first
    _get_executor() creation.
39    #
40    # ? 02 RISK (CODE_CLEANUP_AUDIT S3c): a single worker means ONE stuck or hung job
41    # (e.g. hung Gemini generate, hung WSL GPU call, hung Cloud Run poll) blocks ALL
42    # subsequent jobs ? they stay 'queued' forever. Mitigations already in place:
43    #   - ExecutionPolicy op_timeout_sec kills hung generate/processing calls
44    #   - detector timeout_sec kills hung WSL/GPU calls
45    #   - Cloud Run poll has a max_wait_sec bound
46    #   - ffmpeg/ffprobe subprocesses use bounded media timeouts
47    # A future slice should add a per-job wall-clock timeout (the executor itself
48    # cannot enforce timeouts on the Thread).
49    _job_executor: Optional[ThreadPoolExecutor] = None
50    # Drain worker count, fixed at first _get_executor() creation. 1 (serial) by default;
    configure_drain()
51    # raises it for cloud_run (#466). A module global (not a _get_executor arg) so the many
    call sites +
52    # the tests that monkeypatch `_get_executor` as a no-arg lambda are unchanged.
53    _drain_max_workers: int = 1
54
55
56    def _cloud_dispatch_possible(config: Any = None) -> bool:
57        """True if this server CAN dispatch a job to Cloud Run ? either the profile sets
58        ``deployment.compute_target=cloud_run``, OR the per-request ``compute='cloud'``
    override path is

```

```

59         wired (``CLOUD_RUN_JOB`` + ``GOOGLE_CLOUD_PROJECT`` env + an output bucket). Mirrors
60         orchestration._cloud_available so drain concurrency tracks the ACTUAL remote-
dispatch capability,
61         not just the server default (so compute='cloud' overrides on a local-default server
parallelize too)."""
62         dep = getattr(config, "deployment", None)
63         if getattr(dep, "compute_target", "") == "cloud_run":
64             return True
65         if os.environ.get("CLOUD_RUN_JOB") and os.environ.get("GOOGLE_CLOUD_PROJECT"):
66             stor = getattr(config, "storage", None)
67             if (getattr(stor, "output_root", "") or "").strip():
68                 return True
69         return False
70
71
72     def _resolve_drain_workers(config: Any = None) -> int:
73         """Drain concurrency (#466): a server that CAN dispatch to Cloud Run (see
cloud_dispatch_possible)
74         may run ``deployment.max_concurrent_remote_jobs`` drain workers so I/O-bound remote
dispatch/poll
75         OVERLAPS. This does NOT parallelize CPU/GPU-heavy LOCAL work ? in-process detection
and the compile
76         stitch serialize via orchestration._LOCAL_COMPUTE_LANE regardless of the worker
count ? so the
77         single-box GPU + rate-fragile Gemini invariant holds, and compile/local-override
jobs stay serial.
78         A server with no remote capability stays serial (1). Defaults to 1 when config is
absent
79         (tests / non-server) ? behaviour unchanged."""
80         if _cloud_dispatch_possible(config):
81             dep = getattr(config, "deployment", None)
82             return max(1, int(getattr(dep, "max_concurrent_remote_jobs", 1) or 1))
83         return 1
84
85
86     def configure_drain(config: Any = None) -> None:
87         """Set the drain worker count from ``config`` BEFORE the executor is created (#466).
Call once at
88         server startup. No-op effect once the executor already exists (the count is fixed at
creation)."""
89         global _drain_max_workers
90         _drain_max_workers = _resolve_drain_workers(config)
91
92
93     def _get_executor() -> ThreadPoolExecutor:
94         """Lazy module-global executor; tests monkeypatch this for inline execution. Worker
count is
95         ``_drain_max_workers`` (1 by default; raised for cloud_run via ``configure_drain``,
#466)."""
96         global _job_executor
97         if _job_executor is None:
98             _job_executor = ThreadPoolExecutor(
99                 max_workers=_drain_max_workers, thread_name_prefix="jobworker"
100             )
101         return _job_executor
102
103
104     class JobWaiting(Exception):
105         """Raised inside the pipeline to SELF-SCHEDULE a job (EXECUTION_POLICY.md P3,
WAIT_AND_RESUME). It means "I can't finish now ? park me and resume in `delay_sec``,
106         NOT a failure. The worker persists `next_poll_at` + `progress` (via
`set_job_waiting`),
107         releases the thread, and re-claims the job when due (`claim_due_job`). No work to
date
108         is the segmenter's `execute_processing`, run UNDER the destroy-after-confirm

```



```

166         job,
167         progress=lambda stage: db.set_job_progress(job_id, stage),
168     )
169 else:
170     req = _main._job_to_request(job)
171     result = _main._execute_processing(
172         config,
173         db,
174         req,
175         progress=lambda stage: db.set_job_progress(job_id, stage),
176     )
177     if result.get("video_id"):
178         db.set_job_video_id(job_id, result["video_id"])
179     db.mark_job_done(job_id, result)
180     if not is_compile:
181         _maybe_record_job_cost(
182             job, result, config, db
183         ) # M8 cost ledger (default-OFF; process jobs only)
184         _maybe_run_flywheel(
185             job, result, config
186         ) # M11 data-flywheel (process jobs only)
187 except _main.JobWaiting as w:
188     # Cooperative re-enqueue: persist progress + next_poll_at, release the worker.
189     logger.info(f"Job {job_id} parked for {w.delay_sec:.0f}s (resume when due).")
190     db.set_job_waiting(job_id, w.delay_sec, progress=w.progress)
191 except Exception as e:
192     logger.error(f"Job {job_id} crashed: {e}\n{traceback.format_exc()}")
193     db.mark_job_failed(job_id, f"{type(e).__name__}: {e}")
194
195
196 def _maybe_record_job_cost(
197     job: Dict[str, Any], result: Dict[str, Any], config: Any, db: Any
198 ) -> None:
199     """M8 cost ledger (COMPUTE_DECOUPLED_SERVING M8, D8). DEFAULT-OFF: a no-op unless
200     ``config.billing.enabled``. When on, compute the job's cost from ``result['usage']``
201     (= {gpu_seconds, wall_seconds, egress_gb, gemini_usd, gpu_type}) × the versioned
202     pricebook, and
203     persist it.
204     ? **Honest gap (M8 scope):** the pipeline does **not** emit ``result['usage']`` yet,
205     so even with
206     ``billing.enabled=True`` AND a real calibrated pricebook, every recorded cost is
207     **0** until a
208     FOLLOW-UP wires the worker to emit usage (gpu/wall/egress/gemini). M8 lands the
209     ledger + pricebook
210     + the recording seam; the usage-emission step is separate (the metering hooks on the
211     real run).
212     With the placeholder pricebook the cost is 0 regardless.
213     Best-effort ? never fails a completed job because cost bookkeeping hiccuped. ``job``
214     is the
215     claimed row (carries ``owner_id``); ``result`` is the pipeline output."""
216
217     billing_cfg = getattr(config, "billing", None)
218     if not billing_cfg or not getattr(billing_cfg, "enabled", False):
219         return
220     try:
221         usage = (result or {}).get("usage") or {}
222         db.record_job_cost(
223             job["id"],
224             usage=usage,
225             pricebook_version=billing_cfg.pricebook_version,
226             gpu_type=usage.get("gpu_type"),
227             compute_backend=(result or {}).get("compute_backend"),
228             owner_id=job.get("owner_id"),

```

```

225         margin=billing_cfg.margin,
226     )
227     logger.info(
228         "Recorded cost for job %s (pricebook=%s).",
229         job["id"],
230         billing_cfg.pricebook_version,
231     )
232     except Exception as e: # never wedge a completed job on bookkeeping
233         logger.warning(
234             "Cost recording for job %s failed (non-fatal): %s", job.get("id"), e
235         )
236
237
238     def _maybe_run_flywheel(
239         job: Dict[str, Any], result: Dict[str, Any], config: Any
240     ) -> None:
241         """M11 data-flywheel (COMPUTE_DECOUPLED_SERVING M11, D12). DEFAULT-OFF + INERT: a
no-op unless
242         ``config.flywheel.enabled`` AND attested consent ? and consent is a **faked hard-
deny** today
243         (``flywheel.consent_attested`` returns False for every job until counsel + the
consent schema
244         land), so this captures **nothing** in production. When activated it captures the
consented job's
245         artifacts into the per-video workspace (? later labeling ? the golden training set).
246
247         Best-effort ? never fails a completed job because the flywheel hiccuped. ``job`` is
the claimed
248         row (carries ``owner_id``); ``result`` is the pipeline output."""
249
250         flywheel_cfg = getattr(config, "flywheel", None)
251         if not flywheel_cfg or not getattr(flywheel_cfg, "enabled", False):
252             return
253         try:
254             from backend import flywheel
255
256             flywheel.run_for_job(job, result, config)
257         except Exception as e: # never wedge a completed job on the flywheel
258             logger.warning("Flywheel for job %s failed (non-fatal): %s", job.get("id"), e)
259
260
261     def _drain_due_jobs(config: Any, db: Any) -> int:
262         """One worker tick: claim and run every currently-runnable job (queued, or waiting
whose
263         `next_poll_at` is due), earliest-due first, until the table has nothing runnable.
264
265         Each parked job (`JobWaiting` ? status 'waiting') drops out of this drain and is
picked
266         up by a FUTURE drain once due; we schedule that follow-up drain via
`schedule_drain` so
267         a single-worker box keeps making progress with no central timer (the DB is the
clock).
268         Returns the number of jobs that reached a terminal/parked state this tick.
269         """
270         import backend.main as _main
271
272         ran = 0
273         while True:
274             job = db.claim_due_job()
275             if job is None:
276                 break
277             ran += 1
278             _main._run_claimed_job(job, config, db)
279         # If anything is parked-but-not-yet-due, arrange a wake-up so it resumes without a
new

```

```

280         # submit. On the single-box executor this is a delayed re-drain; multi-worker just
relies
281         # on the next submit/drain. Best-effort ? never raises into the worker.
282         _main._schedule_next_drain_if_waiting(config, db)
283         return ran
284
285
286     def _schedule_next_drain_if_waiting(config: Any, db: Any) -> None:
287         """If any job is parked in 'waiting', arm a re-drain for when the soonest one is
due, so a
288         self-scheduled job resumes with no further client submit. The drain itself re-checks
289         due-ness, so an early wake is harmless and a clamped wake for a far-future job is
fine."""
290         import backend.main as _main
291
292         try:
293             delay = db.seconds_until_next_due()
294         except Exception as e: # pragma: no cover - defensive; never wedge the worker
295             logger.error(f"Could not compute next-due delay: {e}")
296             return
297         if delay is None:
298             return # nothing waiting
299         # Clamp so a far-future job doesn't pin a thread forever; the drain re-checks due-
ness.
300         _main._arm_wake_timer(max(0.0, min(delay, _main._MAX_DRAIN_WAKE_SEC)), config, db)
301
302
303     def _arm_wake_timer(delay_sec: float, config: Any, db: Any) -> None:
304         """Arm a single daemon timer that re-submits a drain after ``delay_sec`` (no central
305         scheduler ? the DB is the clock). Isolated as the sole timer-creating seam so tests
can
306         monkeypatch it to a no-op and assert the *scheduling decision* without real
waiting."""
307         import backend.main as _main
308
309         timer = threading.Timer(
310             delay_sec,
311             lambda: _main._get_executor().submit(_main._drain_due_jobs, config, db),
312         )
313         timer.daemon = True
314         timer.start()
315
316
317     #: A parked job further out than this is re-checked by a capped wake rather than one
long
318     #: sleeping timer ? keeps the wake cadence bounded without a central scheduler.
319     _MAX_DRAIN_WAKE_SEC = 60.0
320
321
322     #
-----
323     # Restart-safe remote jobs (CONTROL_PLANE_DURABLE_REBUILD P3, #471) ? the 2026-07-05
loss shape:
324     # a running cloud detection's done-marker poll lives only in this process's memory, so a
restart
325     # orphaned the job while the L4 worker kept computing; its 29 rallies landed in the
bucket with
326     # no listener. With the md5 persisted at submit (P4), outputs/<md5>/report.json IS a
durable
327     # done-signal the restarted control plane can re-arm on ? completing the job WITHOUT
ever
328     # re-dispatching (no double GPU spend).
329
330     #: Resume-poll cadence. Module-level so tests drive the loop with a zero-wait cadence.
331     _RESUME_POLL EVERY_SEC = 15.0

```

```

332
333
334 def resume_interrupted_remote_jobs(config: Any, db: Any) -> int:
335     """Re-arm the durable done-signal poll for md5-keyed RUNNING process jobs a restart
336     orphaned (the rows ``recover_stale_jobs`` deliberately left alone). Call once at
startup,
337     after ``recover_stale_jobs`` and before any drain. Each job resumes on its own
daemon
338     thread (bounded by the handful of rows a single-instance CP can have in flight);
rows are
339     NEVER re-dispatched from here. Returns the number of resumes started.
340
341     On a box that cannot dispatch to Cloud Run at all (the appliance), no remote worker
can
342     exist ? the rows get today's honest terminal failure instead."""
343     rows = db.interrupted_remote_process_jobs()
344     if not rows:
345         return 0
346     if not _cloud_dispatch_possible(config):
347         for job in rows:
348             db.mark_job_failed(str(job["id"]), "interrupted by server restart")
349             logger.warning(
350                 "[JOBS] %d interrupted md5-keyed process job(s) failed: this server cannot "
351                 "dispatch to Cloud Run, so no remote worker can be in flight.",
352                 len(rows),
353             )
354         return 0
355     # Prove each row was actually CLOUD-dispatched (#494 review round 2): an explicit
356     # compute='cloud' always was; NULL means "server default", which is remote ONLY when
the
357     # default target itself is cloud_run (on a default-local server with cloud-override
358     # capability, a NULL row ran in-process ? no remote worker exists to resume).
359     default_is_cloud = (
360         getattr(getattr(config, "deployment", None), "compute_target", "") ==
"cloud_run"
361     )
362     started = 0
363     for job in rows:
364         # Same normalization the execution path applies (_maybe_dispatch_remote
lowercases
365         # req.compute) ? "LOCAL"/" cloud " must mean at recovery exactly what they meant
at run.
366         requested = (str(job.get("compute") or "")).strip().lower()
367         if requested != "cloud" and not default_is_cloud:
368             db.mark_job_failed(str(job["id"]), "interrupted by server restart")
369             continue
370         t = threading.Thread(
371             target=_resume_one_remote_job,
372             args=(config, db, dict(job)),
373             daemon=True,
374             name=f"resume-{str(job['id'][:8])}",
375         )
376         t.start()
377         started += 1
378     if started:
379         logger.warning(
380             "[JOBS] re-armed the durable done-signal poll for %d interrupted remote
job(s).",
381             started,
382         )
383     return started
384
385
386 def _resume_one_remote_job(config: Any, db: Any, job: Dict[str, Any]) -> None:
387     """One orphaned job's resume: poll ``outputs/<md5>/report.json`` (bounded by the

```

```

same
388     ``deployment.remote_poll_max_wait_sec`` budget a live dispatch gets), then complete
or
389     fail the job HONESTLY from what the bucket says. Never dispatches.""
390     from backend import orchestration, storage
391     from backend.infrastructure.execution_policy import (
392         ExecutionPolicy,
393         PolicyTimeout,
394         poll_until,
395     )
396
397     job_id = str(job["id"])
398     video_id = str(job["video_id"])
399     try:
400         raw_params = job.get("params")
401         params = (
402             raw_params
403             if isinstance(raw_params, dict)
404             else json.loads(raw_params or "{}")
405         )
406         if params.get("force"):
407             # Defense-in-depth (#494 review): the SQL carve-out already excludes forced
rows,
408             # but a forced job must NEVER complete from the bucket cache ? force
deliberately
409             # bypasses it, and nothing proves the report belongs to the forced run
rather
410             # than an older completed one.
411             db.mark_job_failed(
412                 job_id,
413                 "control plane restarted during a forced reprocess; the cached result "
414                 "cannot be trusted for force=true ? re-submit with force",
415             )
416             return
417         if (str(job.get("compute") or "").strip().lower() == "local":
418             # Defense-in-depth (#494 review round 2): an explicit local run has no
remote
419             # worker; a same-md5 bucket report would be an OLDER run's output, not this
job's.
420             db.mark_job_failed(job_id, "interrupted by server restart")
421             return
422         sc = orchestration._storage_settings(config)
423         output_root = ((sc or {}).get("output_root") or "").rstrip("/")
424         if not sc or not output_root:
425             db.mark_job_failed(
426                 job_id,
427                 "interrupted by server restart (no output bucket configured to recover
from)",
428             )
429             return
430         report_uri = f"{output_root}/outputs/{video_id}/report.json"
431         max_wait = float(
432             getattr(getattr(config, "deployment", None), "remote_poll_max_wait_sec",
3900.0)
433             or 3900.0
434         )
435         policy = ExecutionPolicy(
436             poll_every_n_sec=_RESUME_POLL_EVERY_SEC,
437             op_timeout_sec=120.0,
438             max_wait_sec=max_wait,
439         )
440         db.set_job_progress(job_id, "recovering: waiting for the in-flight cloud run")
441         try:
442             poll_until(
443                 lambda: True if storage.exists(report_uri, config=sc) else None,

```

```

444         policy,
445         label=f"resume-{job_id[:8]}",
446         logger=logger,
447         # The #493 live heartbeat: job.progress moves once per still-waiting
tick.
448         on_tick=lambda elapsed: db.set_job_progress(
449             job_id,
450             "recovering: waiting for the in-flight cloud run ? "
451             f"{int(elapsed) // 60} min",
452         ),
453     )
454 except PolicyTimeout:
455     db.mark_job_failed(
456         job_id,
457         "control plane restarted mid-run and the cloud worker's result never "
458         "appeared within the poll budget; re-submit the video (a finished result
"
459         "re-hydrates instantly)",
460     )
461     return
462 payload = orchestration.cached_process_result(
463     db, config, video_id, str(job.get("video_path") or "")
464 )
465 if payload is not None:
466     db.mark_job_done(job_id, payload)
467     logger.warning(
468         "[JOBS] resumed job %s: ingested the in-flight worker's result for %s.",
469         job_id,
470         video_id,
471     )
472     return
473 # The report exists but did not realize into a usable result ? surface its
verdict.
474 status = "unreadable"
475 try:
476     report = orchestration._read_bucket_json(report_uri, sc) or {}
477     status = str(report.get("status") or "unreadable")
478 except Exception: # noqa: BLE001 ? the failure message is best-effort
479     pass
480 db.mark_job_failed(
481     job_id,
482     f"cloud worker finished without a usable result after a control-plane
restart "
483     f"(report status: {status}); re-submit the video",
484 )
485 except Exception as e: # noqa: BLE001 ? a resume crash must land as a job failure,
not vanish
486     logger.error("[JOBS] resume of job %s crashed: %s", job_id, e, exc_info=True)
487     try:
488         db.mark_job_failed(job_id, f"restart recovery failed: {e}")
489     except Exception: # noqa: BLE001
490         pass
491

```

23. user (2026-07-08T05:06:17.321Z)

```

1     import logging
2     import os
3     import subprocess
4     import tempfile
5     from pathlib import Path
6     from typing import Any, Dict, List, Optional, Tuple, Union
7
8     from backend.utils.ffmpeg import ( # P2a: single ffmpeg/ffprobe resolver
9         ffmpeg_bin,

```

```

10         ffmpeg_timeout_sec,
11         ffprobe_bin,
12         ffprobe_timeout_sec,
13         timeout_label,
14     )
15
16     logger = logging.getLogger(__name__)
17
18     # Bundled, commercially-clean default watermark font (Apache-2.0). Shipped so the
overlay
19     # renders out-of-the-box without a system fontconfig setup (a common gap on Windows,
where
20     # `font='Sans'` fails). Overridable via stitching.watermark.font_path. This module lives
at
21     # backend/pipeline/stitcher/, so parents[2] == backend/.
22     _BUNDLED_FONT = (
23         Path(__file__).resolve().parents[2] / "assets" / "fonts" / "RobotoSlab-Regular.ttf"
24     )
25
26
27     class VideoCompiler:
28         def __init__(
29             self,
30             output_dir: str,
31             end_padding: float = 1.0,
32             begin_padding: float = 0.5,
33             watermark: Optional[Any] = None,
34         ):
35             self.output_dir = output_dir
36             self.end_padding = end_padding
37             self.begin_padding = begin_padding
38             # Optional branding overlay. Accepts a WatermarkConfig model or a plain dict;
39             # normalized to a dict (or None) so this module stays config-package-agnostic.
40             self.watermark = self._normalize_watermark(watermark)
41             os.makedirs(output_dir, exist_ok=True)
42
43         @staticmethod
44         def _normalize_watermark(watermark: Optional[Any]) -> Optional[Dict[str, Any]]:
45             if watermark is None:
46                 return None
47             if hasattr(watermark, "model_dump"):
48                 return watermark.model_dump()
49             if isinstance(watermark, dict):
50                 return dict(watermark)
51             return None
52
53         @staticmethod
54         def _ff_quote(path: str) -> str:
55             """Make a filesystem path safe to embed inside single quotes in an ffmpeg
56             filtergraph operand (drawtext fontfile/textfile, movie= source): forward
57             slashes (Windows-friendly), escaped single quotes, AND escaped colons.
58
59             The colon MUST be escaped even inside single quotes: ffmpeg's filtergraph
60             option parser treats ':' as the option separator, so a Windows drive-letter
61             path like ``C:/dir/font.ttf`` otherwise truncates the operand at ``C`` and
62             the encode fails with "Invalid argument". This silently disabled the
63             on-by-default bundled-font watermark (and absolute logo paths) on Windows."""
64             return path.replace("\\", "/").replace("'", r"\'").replace(":", r"\:")
65
66         @staticmethod
67         def _parse_time(val: Union[int, float, str]) -> float:
68             """Normalizes a timestamp value to float seconds.
69             Handles: float/int (pass-through), plain numeric strings ("12.4"),
70             and colon-separated strings ("MM:SS", "HH:MM:SS", "H:MM:SS.ms").
71             This mirrors the parse_time logic used in gemini.py to ensure

```



```

72         the stitcher can always consume whatever format the indexer produced."""
73         if isinstance(val, (int, float)):
74             return float(val)
75         if isinstance(val, str):
76             val = val.strip()
77             if ":" in val:
78                 # Handle MM:SS, HH:MM:SS, etc.
79                 parts = val.split(":")
80                 parts.reverse()
81                 total = 0.0
82                 for i, p in enumerate(parts):
83                     try:
84                         total += float(p) * (60**i)
85                     except ValueError:
86                         pass
87                 return total
88             try:
89                 return float(val)
90             except ValueError:
91                 logger.warning(
92                     f"Could not parse timestamp value '{val}' as a number. Defaulting to
0.0."
93                 )
94                 return 0.0
95         logger.warning(
96             f"Unexpected timestamp type {type(val).__name__} for value '{val}'.
Defaulting to 0.0."
97         )
98         return 0.0
99
100     @staticmethod
101     def _merge_overlapping(
102         segments: List[Tuple[float, float]],
103     ) -> List[Tuple[float, float]]:
104         """Normalize, sort, and merge overlapping/duplicate (start, end) ranges.
105
106         A highlight reel must never replay the same footage. The same rally can land in
107         the segment list more than once ? e.g. two detector runs accumulated for one
108         video (add_video is an UPSERT, so re-running without a prune leaves both sets),
109         or
110         chunk-overlap duplicates near boundaries. Stitching those verbatim plays the
111         rally
112         twice. Merging any range that overlaps or abuts the previous one collapses true
113         duplicates/overlaps while preserving genuinely distinct (gapped) rallies.
114         Degenerate rows (end <= start) are dropped. Timestamps may be float/int/str.
115         """
116         parsed = []
117         for raw_s, raw_e in segments:
118             s = VideoCompiler._parse_time(raw_s)
119             e = VideoCompiler._parse_time(raw_e)
120             if e > s:
121                 parsed.append((s, e))
122         parsed.sort()
123         merged: List[Tuple[float, float]] = []
124         for s, e in parsed:
125             if merged and s <= merged[-1][1]: # overlaps/abuts previous -> extend it
126                 merged[-1] = (merged[-1][0], max(merged[-1][1], e))
127             else:
128                 merged.append((s, e))
129         return merged
130
131     def _get_video_duration(self, video_path: str) -> float:
132         """Probes the video file to get its total duration in seconds.
Returns 0.0 if it cannot be determined."""
try:

```

```

133         timeout = ffprobe_timeout_sec()
134         result = subprocess.run(
135             [
136                 ffprobe_bin(),
137                 "-v",
138                 "error",
139                 "-show_entries",
140                 "format=duration",
141                 "-of",
142                 "default=noprint_wrappers=1:nokey=1",
143                 video_path,
144             ],
145             capture_output=True,
146             text=True,
147             check=True,
148             timeout=timeout,
149         )
150         return float(result.stdout.strip())
151     except subprocess.TimeoutExpired:
152         logger.warning(
153             "Could not determine video duration via ffprobe: timed out after %s",
154             timeout_label(ffprobe_timeout_sec()),
155         )
156         return 0.0
157     except Exception as e:
158         logger.warning(f"Could not determine video duration via ffprobe: {e}")
159         return 0.0
160
161 def _watermark_filter(self, tmp_dir: str) -> Optional[str]:
162     """Build the ffmpeg -vf filtergraph that burns the configured watermark (text
163     and/or logo PNG) into a clip during the per-clip re-encode. Returns None when
164     watermarking is disabled or has nothing to draw. The concat stage stays -c copy
165     because the mark is already baked into every clip identically."""
166     wm = self.watermark
167     if not wm or not wm.get("enabled"):
168         return None
169
170     text = (wm.get("text") or "").strip()
171     logo_path = (wm.get("logo_path") or "").strip()
172     if logo_path and not os.path.exists(logo_path):
173         logger.warning(
174             f"[watermark] logo not found: {logo_path} ? skipping logo overlay."
175         )
176         logo_path = ""
177     if not text and not logo_path:
178         return None
179
180     margin = int(wm.get("margin", 24))
181     opacity = float(wm.get("opacity", 0.85))
182     position = str(wm.get("position", "bottom-right"))
183     # drawtext uses lowercase w/h (main) + tw/th (text); overlay uses W/H (main) +
w/h (logo).
184     dt_xy = {
185         "bottom-right": f"x=w-tw-{margin}:y=h-th-{margin}",
186         "bottom-left": f"x={margin}:y=h-th-{margin}",
187         "top-right": f"x=w-tw-{margin}:y={margin}",
188         "top-left": f"x={margin}:y={margin}",
189     }.get(position, f"x=w-tw-{margin}:y=h-th-{margin}")
190     ov_xy = {
191         "bottom-right": f"W-w-{margin}:H-h-{margin}",
192         "bottom-left": f"{margin}:H-h-{margin}",
193         "top-right": f"W-w-{margin}:{margin}",
194         "top-left": f"{margin}:{margin}",
195     }.get(position, f"W-w-{margin}:H-h-{margin}")
196

```

```

197         drawtext = ""
198         if text:
199             ratio = float(wm.get("font_size_ratio") or 0.035)
200             divisor = max(1, round(1.0 / ratio))
201             # textfile= sidesteps the notorious drawtext text= escaping for arbitrary
strings.
202             tf_path = os.path.join(tmp_dir, "watermark.txt")
203             with open(tf_path, "w", encoding="utf-8") as fh:
204                 fh.write(text)
205             # Font precedence: explicit font_path ? bundled Apache-2.0 default (renders
206             # without fontconfig) ? fontconfig family (last-resort, environment-
dependent).
207             font_path = (wm.get("font_path") or "").strip()
208             if not font_path and _BUNDLED_FONT.exists():
209                 font_path = str(_BUNDLED_FONT)
210             if font_path:
211                 font_opt = f"fontfile='{self._ff_quote(font_path)}'"
212             else:
213                 font_opt = f"font='{wm.get('font', 'Sans')}'"
214             box = ""
215             if wm.get("box", True):
216                 box = f":box=1:boxcolor=black@{min(opacity, 0.5):.2f}:boxborderw=10"
217             drawtext = (
218                 f"drawtext=textfile='{self._ff_quote(tf_path)}':{font_opt}"
219                 f":fontcolor=white@{opacity:.2f}:fontsize=h/{divisor}:{dt_xy}{box}"
220             )
221
222         if logo_path and drawtext:
223             return
224         f"movie='{self._ff_quote(logo_path)}'[wm];[in][wm]overlay={ov_xy},{drawtext}"
225         if logo_path:
226             return f"movie='{self._ff_quote(logo_path)}'[wm];[in][wm]overlay={ov_xy}"
227         return drawtext
228
229     def _trim_one_segment(
230         self,
231         idx: int,
232         raw_start: Union[int, float, str],
233         raw_end: Union[int, float, str],
234         segments: List[Tuple[float, float]],
235         video_duration: float,
236         tmp_dir: str,
237         raw_video_path: str,
238         watermark_filter: Optional[str],
239     ) -> Tuple[Optional[str], Optional[Tuple[int, float, float, str]]]:
240         """Trim a single segment into its own clip file (extracted verbatim from the
241         per-segment body of compile_highlights). Returns (clip_path, skip_record): on
242         success the clip path and None; on a skip (invalid/out-of-range/ffmpeg failure)
243         None and the (idx, start, end, reason) record to record in skipped_segments."""
244         # Normalize timestamps to float seconds ? handles float, int,
245         # plain numeric strings, and colon-separated formats like "MM:SS"
246         start = self._parse_time(raw_start)
247         end = self._parse_time(raw_end)
248         segment_label = (
249             f"Segment #{idx + 1}/{len(segments)} [{start:.2f}s - {end:.2f}s]"
250         )
251
252         # Validate the segment times
253         if start < 0:
254             logger.warning(
255                 f"{segment_label}: start_time is negative ({start:.2f}s). Clamping to
0."
256             )
257             start = 0.0

```

```

258         if start >= end:
259             logger.error(
260                 f"{segment_label}: start_time ({start:.2f}s) >= end_time ({end:.2f}s).
Skipping invalid segment."
261             )
262             return None, (idx, start, end, "start >= end")
263
264         # Check if segment extends beyond video duration
265         if video_duration > 0:
266             if start >= video_duration:
267                 logger.error(
268                     f"{segment_label}: start_time ({start:.2f}s) is beyond the video
duration ({video_duration:.2f}s). "
269                     f"This segment cannot be extracted ? the video ran out. Skipping."
270                 )
271                 return None, (idx, start, end, "start beyond video end")
272
273             if end > video_duration:
274                 original_end = end
275                 end = video_duration
276                 logger.warning(
277                     f"{segment_label}: end_time ({original_end:.2f}s) exceeds video
duration ({video_duration:.2f}s). "
278                     f"Clamping end_time to {video_duration:.2f}s. The segment may be
truncated."
279                 )
280
281             clip_path = os.path.join(tmp_dir, f"clip_{idx}.mp4")
282
283             # Precise sub-second trim using fast re-encoding to preserve stream headers
284             padded_start = max(0.0, start - self.begin_padding)
285             padded_end = end + self.end_padding
286
287             # Clamp padded_end to video duration if known
288             if video_duration > 0 and padded_end > video_duration:
289                 padded_end = video_duration
290                 logger.info(
291                     f"{segment_label}: Padded end ({end + self.end_padding:.2f}s) exceeds
video duration. "
292                     f"Clamping to {video_duration:.2f}s (end padding partially applied)."
293                 )
294
295             trim_duration = padded_end - padded_start
296             logger.info(
297                 f"Trimming {segment_label} -> [{padded_start:.3f}s - {padded_end:.3f}s]
(duration: {trim_duration:.2f}s)"
298             )
299
300             base_cmd = [
301                 ffmpeg_bin(),
302                 "-y",
303                 "-ss",
304                 str(round(padded_start, 3)),
305                 "-to",
306                 str(round(padded_end, 3)),
307                 "-i",
308                 raw_video_path,
309                 "-c:v",
310                 "libx264",
311                 "-preset",
312                 "superfast",
313                 "-crf",
314                 "22",
315                 "-c:a",
316                 "aac",

```

```

317         "-b:a",
318         "128k",
319     ]
320     vf_args = ["-vf", watermark_filter] if watermark_filter else []
321
322     # Try with the watermark first; if that encode fails (e.g. a bad font/logo
323     # path), retry once WITHOUT it so a branding misconfig can never nuke the
324     # whole reel. Clips stay codec-identical either way, so concat -c copy holds.
325     attempts = [vf_args, []] if vf_args else [[]]
326     produced = False
327     last_exc: Optional[subprocess.CalledProcessError] = None
328     cmd = base_cmd + [clip_path]
329     timeout = ffmpeg_timeout_sec()
330     for attempt_vf in attempts:
331         cmd = base_cmd + attempt_vf + [clip_path]
332         try:
333             subprocess.run(
334                 cmd, capture_output=True, check=True, timeout=timeout
335             )
336         except subprocess.TimeoutExpired:
337             reason = f"ffmpeg timeout after {timeout_label(timeout)}"
338             logger.error(f"{segment_label}: FFmpeg trim timed out.")
339             print(f"  Command: {' '.join(cmd)}")
340             print(f"  {reason}")
341             return None, (idx, start, end, reason)
342         except subprocess.CalledProcessError as e:
343             last_exc = e
344             if attempt_vf and vf_args:
345                 _err = (
346                     e.stderr.decode(errors="replace").strip()[-300:]
347                     if getattr(e, "stderr", None)
348                     else ""
349                 )
350                 logger.warning(
351                     f"{segment_label}: watermark encode failed; retrying without it."
352                     f"ffmpeg: {_err or '(no stderr captured)'}"
353                 )
354                 continue
355             break
356     if os.path.exists(clip_path) and os.path.getsize(clip_path) > 0:
357         produced = True
358         if not attempt_vf and vf_args:
359             logger.warning(
360                 f"{segment_label}: encoded WITHOUT the watermark (the watermark
filter failed)."
361             )
362         break
363     last_exc = None # ffmpeg succeeded but produced nothing
364     break
365
366     if produced:
367         return clip_path, None
368     elif last_exc is not None:
369         stderr_msg = last_exc.stderr.decode(errors="replace").strip()
370         logger.error(f"{segment_label}: FFmpeg trim failed.")
371         print(f"  Command: {' '.join(cmd)}")
372         print(
373             f"  FFmpeg stderr: {stderr_msg[-500:]}")
374         ) # Last 500 chars to avoid flooding
375         return None, (idx, start, end, f"ffmpeg error: {stderr_msg[-200:]}")
376     else:
377         logger.error(
378             f"{segment_label}: FFmpeg exited successfully but output clip is missing
or empty. Skipping."

```

```

379         )
380         return None, (idx, start, end, "empty output clip")
381
382     def compile_highlights(
383         self,
384         raw_video_path: str,
385         segments: List[Tuple[float, float]],
386         output_filename: str,
387     ) -> str:
388         """
389         Extracts selected segments from a raw video and stitches them together
390         into a seamless highlights file, applying end_padding to prevent abrupt endings.
391
392         .. warning::
393             This method uses a two-stage codec contract: per-clip re-encodes must use
394             IDENTICAL encoder settings (codec, preset, CRF, pixel format) so the final
395             concat can stream-copy (``-c copy``) without re-encoding. Changing the
396             clip encode settings without updating the concat assumptions will produce
397             broken output.
398         """
399         if not segments:
400             raise ValueError("No highlight segments provided to compile.")
401
402         # Never stitch the same footage twice: collapse overlapping/duplicate ranges
403         # (e.g. accumulated re-run segments or chunk-overlap dupes) before trimming.
404         original_count = len(segments)
405         segments = self._merge_overlapping(segments)
406         if len(segments) < original_count:
407             logger.info(
408                 "Merged %d overlapping/duplicate segment(s) before stitching: %d ->
%d.",
409                 original_count - len(segments),
410                 original_count,
411                 len(segments),
412             )
413         if not segments:
414             raise ValueError("No valid (start < end) highlight segments after merge.")
415
416         if not os.path.exists(raw_video_path):
417             raise FileNotFoundError(
418                 f"[COMPILER ERROR] Source video file not found: {raw_video_path}"
419             )
420
421         # Defense-in-depth: never let the output name escape output_dir (the API also
422         # validates this at the request boundary).
423         output_path = os.path.join(self.output_dir, os.path.basename(output_filename))
424
425         # Probe video duration so we can detect out-of-range segments
426         video_duration = self._get_video_duration(raw_video_path)
427         if video_duration > 0:
428             logger.info(f"Source video duration: {video_duration:.2f}s")
429         else:
430             logger.warning(
431                 "Video duration unknown. Cannot validate segment boundaries against
video length."
432             )
433
434         # We will create temporary clips and stitch them using FFmpeg concat demuxer
435         temp_clips = []
436         skipped_segments = []
437
438         # Create a temp directory within output_dir
439         with tempfile.TemporaryDirectory(dir=self.output_dir) as tmp_dir:
440             logger.info(
441                 f"Trimming {len(segments)} segments

```

```

442         )
443         # Build the (optional) watermark filtergraph once; it is applied during
every
444         # per-clip re-encode so the mark is baked into each clip identically.
445         watermark_filter = self._watermark_filter(tmp_dir)
446         if watermark_filter:
447             logger.info("[watermark] applying branding overlay to each clip.")
448         for idx, (raw_start, raw_end) in enumerate(segments):
449             clip_path, skip_record = self._trim_one_segment(
450                 idx,
451                 raw_start,
452                 raw_end,
453                 segments,
454                 video_duration,
455                 tmp_dir,
456                 raw_video_path,
457                 watermark_filter,
458             )
459             if clip_path is not None:
460                 temp_clips.append(clip_path)
461             if skip_record is not None:
462                 skipped_segments.append(skip_record)
463
464         # Report summary of skipped segments
465         if skipped_segments:
466             logger.info(
467                 f"\n[SUMMARY] {len(skipped_segments)} out of {len(segments)}
segments were SKIPPED:"
468             )
469             for seg_idx, seg_start, seg_end, reason in skipped_segments:
470                 logger.info(
471                     f" - Segment #{seg_idx + 1} [{seg_start:.2f}s -
{seg_end:.2f}s]: {reason}"
472                 )
473             logger.info("")
474
475         if not temp_clips:
476             raise RuntimeError(
477                 f"[COMPILER ERROR] Failed to generate ANY valid highlight clips. "
478                 f"All {len(segments)} segments were skipped. "
479                 f"Check the errors above for details. "
480                 f"Common causes: video file is corrupted, segments are out of range,
or ffmpeg is not installed."
481             )
482
483         logger.info(
484             f"Successfully trimmed {len(temp_clips)}/{len(segments)} clips.
Preparing to stitch..."
485         )
486
487         # Create a list file for FFmpeg concat demuxer
488         list_file_path = os.path.join(tmp_dir, "clips.txt")
489         # encoding pinned: a non-ASCII clip path would crash on Windows (cp1252
default);
490         # ffmpeg's concat demuxer reads the list as UTF-8 on every OS.
491         with open(list_file_path, "w", encoding="utf-8") as f:
492             for clip in temp_clips:
493                 # The concat demuxer resolves relative entries against the list
file's
494                 # directory. Use absolute paths so a temp-dir clip does not become
495                 # output/tmp.../output/tmp.../clip_0.mp4 on Windows.
496                 safe_path = Path(clip).resolve().as_posix().replace("'", "'\\'")
497                 f.write(f"file '{safe_path}'\n")
498

```

```

499         logger.info("Stitching clips together...")
500         # Stitch them together without re-encoding (since they were encoded with
identical settings)
501         concat_cmd = [
502             ffmpeg_bin(),
503             "-y",
504             "-f",
505             "concat",
506             "-safe",
507             "0",
508             "-i",
509             list_file_path,
510             "-c",
511             "copy",
512             output_path,
513         ]
514
515         try:
516             timeout = ffmpeg_timeout_sec()
517             subprocess.run(
518                 concat_cmd, capture_output=True, check=True, timeout=timeout
519             )
520
521             # Final verification
522             if not os.path.exists(output_path) or os.path.getsize(output_path) == 0:
523                 raise RuntimeError(
524                     f"[COMPILER ERROR] FFmpeg concatenation appeared to succeed, "
525                     f"but output file is missing or empty: {output_path}"
526                 )
527
528             output_size_mb = os.path.getsize(output_path) / (1024 * 1024)
529             logger.info(
530                 f"Dynamic highlight reel saved: {output_path} ({output_size_mb:.1f}
MB)"
531             )
532             return output_path
533         except subprocess.CalledProcessError as e:
534             stderr_msg = e.stderr.decode(errors="replace").strip()
535             raise RuntimeError(
536                 f"[COMPILER ERROR] FFmpeg concatenation failed.\n"
537                 f"  Command: {' '.join(concat_cmd)}\n"
538                 f"  FFmpeg stderr: {stderr_msg[-500:]}"
539             ) from e
540         except subprocess.TimeoutExpired as e:
541             raise RuntimeError(
542                 f"[COMPILER ERROR] FFmpeg concatenation timed out after
{timeout_label(timeout)}.\n"
543                 f"  Command: {' '.join(concat_cmd)}"
544             ) from e
545

```

24. assistant (2026-07-08T05:06:21.405Z)

Now let me find the compile API endpoint, the job submission, and the cloud-serving preset.

25. user (2026-07-08T05:06:24.016Z)

```

67:from backend.pipeline.stitcher.compiler import VideoCompiler
228:# The processing/compile ORCHESTRATION (hash?validate?segment?report pipeline, Cloud Run
229:# dispatch+ingest, compile resolution + worker handler) now lives in backend.orchestration
232:# call-time `_main._execute_compile(...)` seam, and monkeypatch.setattr(backend.main, ...)
237:     _execute_compile,
244:     _resolve_compile,
479:     """List compiled highlight reels present in the output dir, newest first, with a
download URL.

```



```

481:     NB: reels are stored FLAT in output_dir under the bare name passed to /api/compile ? the
engine
483:     this is a global list. (Per-video reel tracking would need a DB record + compile change
? a
512:     """Stream a compiled reel by its bare filename (no video_id ? reels are flat in
output_dir).
525:         raise HTTPException(status_code=404, detail=f"No compiled reel named '{name}'.")
1043:     """Stream a compiled highlight reel ? `filename` is the bare name given to /api/compile
1076:         detail=f"No compiled file named '{name}'. Compile first via /api/compile.",
1098:@app.post("/api/compile", status_code=202)
1099:def compile_highlights(req: CompileRequest, request: Request):
1114:     video, _kept, out_name = _resolve_compile(
1127:     if not request.app.state.db.create_compile_job(
1137:         status_code=500, detail="Failed to persist compile job record."
1381:     compiler = VideoCompiler(
1388:         res_path = compiler.compile_highlights(
1393:         print(f"[ERROR] Failed to compile trimmed segment: {e}")
1429:     # /api/compile ? reads the same value instead of a write-only runtime hack.
1480:         f"[JOBS] Recovered stale jobs from a previous run ? {requeued} compile re-
queued, "

```

26. user (2026-07-08T05:06:24.395Z)

```

1     """Deployment-profile presets + environment auto-detection (COMPUTE_DECOUPLED_SERVING
M1).
2
3     A *deployment profile* is a single switch that applies a coherent BUNDLE of existing
4     config knobs, so the truly-local / cloud-serving / cpu-mock modes don't drift apart one
5     knob at a time (the design in ``docs/COMPUTE_DECOUPLED_SERVING/README.md`` §4.3).
6
7     Precedence (lowest ? highest), realised by the loader::
8
9         built-in model defaults < profile preset < config.json < env < worker --set
10
11     This module is **pure data + pure functions** ? no IO, and (critically) **no torch
import at
12     module load**. The loader applies a preset ONLY when a profile is EXPLICITLY selected
13     (``config.json`` ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); with no
14     profile selected the loader is byte-identical to the pre-M1 behaviour (``profile=""`` ==
15     today). Auto-detect only ever **suggests** a profile ? it never auto-applies one (D15:
cloud
16     spend is never entered silently).
17     """
18
19     from __future__ import annotations
20
21     import copy
22     import os
23     from typing import Any, Mapping, Optional
24
25     # Canonical enum values, defined ONCE here. ``models.DeploymentConfig``'s ``Literal``
mirrors
26     # these and ``tests/test_config_deployment.py`` pins parity so the two can never drift.
27     PROFILES: tuple[str, ...] = ("truly-local", "cloud-serving", "cpu-mock")
28     COMPUTE_TARGETS: tuple[str, ...] = ("local_gpu", "cloud_run", "cpu_mock")
29
30     # Each profile maps 1:1 to its canonical compute_target. Used to apply the preset and to
warn
31     # when config.json explicitly overrides compute_target into a state inconsistent with
the profile.
32     COMPUTE_TARGET_FOR_PROFILE: dict[str, str] = {
33         "truly-local": "local_gpu",
34         "cloud-serving": "cloud_run",
35         "cpu-mock": "cpu_mock",
36     }

```

```

37
38     # A preset is a partial, nested override dict over the model defaults. It touches ONLY
knobs
39     # that EXIST today ? M1 is pure/additive: the new
``deployment.{profile,compute_target}`` plus
40     # ``indexing.{detector_impl,skip_ai_handoff}`` and ``storage.backend``. Later milestones
extend
41     # these bundles ? ``input_backend`` (M2) and the configurable output/workspace base (M3,
42     # Launch-Report-gated) ? deliberately NOT set here so M1 never pre-empts a gated flip.
43     PROFILE_PRESETS: dict[str, dict[str, Any]] = {
44         "truly-local": {
45             "deployment": {"profile": "truly-local", "compute_target": "local_gpu"},
46             # detector_impl stays "auto" (the runner picks WSL vs native by env on the GPU
box).
47             # Zero-cloud / privacy-first ? run fully local with no Gemini handoff.
48             "indexing": {"skip_ai_handoff": True},
49             "storage": {"backend": "local"},
50         },
51         "cloud-serving": {
52             "deployment": {"profile": "cloud-serving", "compute_target": "cloud_run"},
53             # Linux L4 (cuda); Gemini handoff ON until the owned model clears the #172 floor
(D11).
54             # wasb GPU-feed tuning: the parity defaults (batch 8, no prefetch) left the paid
L4
55             # <1% utilized on the first real serving run (2026-07-03) ? the pipeline was
56             # CPU-feed-bound. Serving opts into the measured fast values; local/WSL
unchanged.
57             "indexing": {
58                 "detector_impl": "native",
59                 "skip_ai_handoff": False,
60                 # GPU-resident NVDEC decode (#506/#507): validated on a real L4 (golden gate
PASS,
61                 # F1 0.574 ? baseline 0.582 @ SERVED) AND on the live Cloud Run L4 (the
served
62                 # trajectory reproduced the VM run ? 26479 ? 26481 visible on GX010128),
then rolled
63                 # out + canaried (DEPLOY_REPORT_cloud-serving_gpu-resident-
nvdec_2026-07-07). This makes
64                 # nvdec the code-level serving default so it survives an image rebuild; the
worker's
65                 # WASB_DECODE_BACKEND env was the reversible rollout lever. cuda-only
(validator-enforced).
66                 "wasb": {
67                     "batch_size": 64,
68                     "prefetch_batches": 4,
69                     "decode_backend": "nvdec_resident",
70                 },
71             },
72             "storage": {"backend": "gcs"},
73         },
74         "cpu-mock": {
75             "deployment": {"profile": "cpu-mock", "compute_target": "cpu_mock"},
76             # Deterministic CPU mock ? no GPU/WSL, no provider/API. The CI / regression
profile.
77             "indexing": {"detector_impl": "stub", "skip_ai_handoff": True},
78             "storage": {"backend": "local"},
79         },
80     }
81
82
83     def is_known_profile(profile: str) -> bool:
84         """True for a non-empty, recognised profile name."""
85         return bool(profile) and profile in PROFILE_PRESETS
86
87

```

```

88     def preset_for(profile: str) -> dict[str, Any]:
89         """Return a deep COPY of the preset bundle for ``profile`` (so callers can't mutate
the
90         module-level table), or ``{}`` for ``""`` / an unknown profile."""
91         return copy.deepcopy(PROFILE_PRESETS.get(profile, {}))
92
93
94     def _truthy(val: Optional[str]) -> bool:
95         """Common env-var truthiness (``1/true/yes/on``, case/space-insensitive)."""
96         if val is None:
97             return False
98         return val.strip().lower() in {"1", "true", "yes", "on"}
99
100
101     def _cuda_available() -> bool:
102         """Best-effort CUDA probe. torch is heavy + optional, so it is imported lazily here
and
103         ONLY when a caller opts into ``probe_cuda``; any import/runtime failure means 'no
GPU'."""
104         try:
105             # Intentionally lazy + local: keeps the config-load path torch-free (torch is
heavy
106             # and optional). Only reached when a caller opts into probe_cuda.
107             import torch
108
109             return bool(torch.cuda.is_available())
110         except Exception:
111             return False
112
113
114     def probe_cuda_available() -> bool:
115         """Public best-effort CUDA probe (lazy torch import; ``False`` on any failure).
116
117         Exposed for the per-profile startup checks (``backend/config/startup.py``) so a
caller
118         that *wants* the GPU heads-up can pass the result in without the startup module
importing
119         torch itself. Best-effort by design: if torch lives only in a detector subprocess
env
120         (native WASB), this returns ``False`` in the main process ? which is why the startup
checks
121         treat a GPU mismatch as a WARNING (the runner healthcheck is the authority), never a
hard
122         error that could false-fail."""
123         return _cuda_available()
124
125
126     def suggest_profile(
127         env: Optional[Mapping[str, str]] = None, *, probe_cuda: bool = False
128     ) -> str:
129         """SUGGEST (never auto-apply ? D15) a deployment profile from the environment.
130
131         Signals, in priority order:
132
133         - ``K_SERVICE`` set (Cloud Run)           ? ``"cloud-serving"``
134         - ``CI`` truthy                            ? ``"cpu-mock"``
135         - ``probe_cuda`` and a CUDA GPU present ? ``"truly-local"``
136         - otherwise                               ? ``""`` (no strong suggestion)
137
138         ``probe_cuda`` defaults to ``False`` so the common config-load path stays cheap and
never
139         imports torch; a CLI ``--detect`` or cloud-serving startup that wants the GPU probe
passes
140         ``probe_cuda=True``. The result is advisory only: the caller decides whether to act
on it.

```

```

141         """
142         e = os.environ if env is None else env
143         if e.get("K_SERVICE"):
144             return "cloud-serving"
145         if _truthy(e.get("CI")):
146             return "cpu-mock"
147         if probe_cuda and _cuda_available():
148             return "truly-local"
149         return ""
150

```

27. assistant (2026-07-08T05:06:27.370Z)

Let me read the compile API endpoint and how compile jobs store their params.

28. user (2026-07-08T05:06:29.382Z)

```

1080
1081
1082 @app.post("/api/query")
1083 def query_play_segments(req: QueryRequest, request: Request):
1084     filters = {
1085         "server": req.server,
1086         "receiver": req.receiver,
1087         "duration_min": req.duration_min,
1088         "event_type": req.event_type,
1089     }
1090     segments = request.app.state.db.query_play_segments(req.video_id, filters)
1091     if not segments and read_path_rehydrate(
1092         request.app.state.db, request.app.state.config, req.video_id
1093     ): # #485: DB = cache, bucket = truth ? re-hydrate before reporting nothing
1094         segments = request.app.state.db.query_play_segments(req.video_id, filters)
1095     return {"play_segments": segments}
1096
1097
1098 @app.post("/api/compile", status_code=202)
1099 def compile_highlights(req: CompileRequest, request: Request):
1100     """Submit a reel-COMPILE job (202 + job_id); poll GET /api/jobs/{job_id} for the
1101     reel (#329).
1102     Async, mirroring /api/process: the ffmpeg stitch (trim + re-encode each clip +
1103     concat + watermark)
1104     can run minutes on CPU ? past the ~100 s edge/proxy timeout ? so doing it
1105     synchronously returned a
1106     Cloudflare 524 while the engine kept going (the reel built, but the client saw
1107     "failed" and never
1108     got the URL). The cheap checks (video exists, segments match + survive the floor,
1109     safe filename)
1110     fast-fail HERE with a 4xx; the stitch runs on the worker. On `done`, the job's
1111     `result` carries
1112     {filepath, download_url} ? the same shape the sync endpoint returned, now under
1113     /api/jobs/{id}."""
1114     # M9 edge-auth tenant (DEFAULT-OFF ? None); a missing/spoofed identity fails 401
1115     before queuing.
1116     try:
1117         owner_id = edge_auth.resolve_tenant(request.headers, request.app.state.config)
1118     except edge_auth.EdgeAuthError as e:
1119         raise HTTPException(status_code=401, detail=f"edge auth: {e}") from e
1120     try:
1121         video, _kept, out_name = _resolve_compile(
1122             request.app.state.db,
1123             request.app.state.config,
1124             req.video_id,
1125             req.segment_ids,
1126             req.output_filename,

```

```

1120         )
1121     except _CompileError as e:
1122         raise HTTPException(status_code=e.status_code, detail=e.detail) from e
1123
1124     job_id = uuid.uuid4().hex
1125     # Persist the video's SOURCE ref (the NOT NULL video_path); the worker re-fetches
the row by id so
1126     # the freshest filepath wins. owner_id/locale ride along for M9 multi-tenant + the
localized reel.
1127     if not request.app.state.db.create_compile_job(
1128         job_id,
1129         req.video_id,
1130         video["filepath"],
1131         req.segment_ids,
1132         out_name,
1133         locale=req.locale,
1134         owner_id=owner_id,
1135     ):
1136         raise HTTPException(
1137             status_code=500, detail="Failed to persist compile job record."
1138         )
1139     _get_executor().submit(
1140         _drain_due_jobs, request.app.state.config, request.app.state.db
1141     )
1142     return JSONResponse(
1143         status_code=202,
1144         content={"job_id": job_id, "status": "queued", "poll": f"/api/jobs/{job_id}"},
1145     )
1146
1147
1148 # --- CLI Debug Utility & Orchestration ---
1149 def run_cli_diagnose_clip(video_path: str, timestamp: float):
1150     """CLI debugging utility to examine segment properties around a timestamp."""
1151     print("=" * 60)
1152     print(f"DEBUGGING VIDEO CLIP: {video_path} at {timestamp}s")
1153     print("=" * 60)
1154
1155     cfg = load_config() # legacy wrapper still works, returns dict for now
1156
1157     # 1. Validation check diagnostics
1158     print("[DIAGNOSTIC] Running validation framework...")
1159     results = registry.run_all(video_path, cfg)
1160     for r in results:
1161         status_symbol = "[PASS]" if r.passed else "[FAIL]"
1162         print(f" {status_symbol} {r.validator_name}: {r.message}")
1163         if r.details:
1164             print(f"      Details: {r.details}")
1165
1166     # 2. Local segment analysis around the timestamp (sampling +/- 5 seconds)
1167     print("-" * 60)
1168     print("[DIAGNOSTIC] Analyzing local motion around the timestamp...")
1169     # Lazy import: cv2 is NOT installed in the LIGHT tier (torch-free default)
1170     # and this --debug-clip path is CLI-only, not on the served API path.
1171     import cv2
1172
1173     cap = cv2.VideoCapture(video_path)
1174     if not cap.isOpened():
1175         print("[FAIL] OpenCV could not open video.")
1176         return
1177
1178     fps = cap.get(cv2.CAP_PROP_FPS)
1179     total_frames = cap.get(cv2.CAP_PROP_FRAME_COUNT)
1180
1181     start_frame = max(0, int((timestamp - 3.0) * fps))
1182     end_frame = min(int(total_frames), int((timestamp + 3.0) * fps))

```

```

1183
1184     # Measure frame-by-frame changes in sample region
1185     cap.set(cv2.CAP_PROP_POS_FRAMES, start_frame)
1186     prev_gray = None
1187     motions = []
1188
1189     for _ in range(start_frame, end_frame):
1190         ret, frame = cap.read()
1191         if not ret:
1192             break
1193         gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
1194         gray = cv2.GaussianBlur(gray, (21, 21), 0)
1195
1196         if prev_gray is not None:
1197             diff = cv2.absdiff(prev_gray, gray)
1198             _, thresh = cv2.threshold(diff, 25, 255, cv2.THRESH_BINARY)
1199             motion_ratio = cv2.countNonZero(thresh) / (
1200                 thresh.shape[0] * thresh.shape[1]
1201             )
1202             motions.append(motion_ratio)
1203             prev_gray = gray
1204
1205     cap.release()
1206
1207     if motions:
1208         avg_motion = sum(motions) / len(motions)
1209         # Support both legacy dict (from wrapper) and future direct model
1210         if hasattr(cfg, "indexing"):
1211             threshold = getattr(cfg.indexing, "motion_threshold", 0.08)
1212         else:
1213             threshold = cfg.get("indexing", {}).get("motion_threshold", 0.08) # type:
ignore[attr-defined]
1214         print(f" Sample Frame Count: {len(motions)}")
1215         print(
1216             f" Average motion index around timestamp: {round(avg_motion, 4)}
(threshold: {threshold})"
1217         )
1218         if avg_motion > threshold:
1219             print(
1220                 " Result: [ACTIVE PLAY] High motion detected. Segment classifies as
active play."
1221             )
1222         else:
1223             print(
1224                 " Result: [DEAD TIME] Idle movement. Segment classified as dead time."
1225             )
1226     else:
1227         print(" [ERROR] No visual frames were extracted in the sample range.")
1228
1229     print("=" * 60)
1230
1231
1232     def _enforce_startup_or_exit(config) -> None:
1233         """M5: run the startup checks for the server path and ``sys.exit(1)`` on a FATAL
finding BEFORE
1234         any work ? a fatal incoherent profile (e.g. cloud-serving pointed at local-only
storage) OR, for
1235         the exposed server (``include_exposure=True``, M9 hardening), a half-configured
exposure (edge
1236         auth on without the ``allowed_video_root`` jail / CF team+AUD). A strict no-op when
no
1237         ``deployment.profile`` is selected AND edge auth is off (today's behaviour). No GPU
probe here ?
1238         the server stays torch-light; the detector healthcheck is the GPU authority.
Extracted so the

```

1239 fail-fast wiring + the default-OFF no-op are unit-testable."""

29. user (2026-07-08T05:06:30.709Z)

```
462:      # v9: the jobs table now holds TWO kinds of work ? `kind`='process' (the original,
NULL ?
463:      # process for pre-v9 rows) and `kind`='compile' (the async reel stitch, #329;
/api/compile
465:      # is the kind-specific JSON payload a compile job carries
(segment_ids/output_filename/locale)
467:      self._add_column_idempotent(cursor, "ALTER TABLE jobs ADD COLUMN kind TEXT")
```

30. assistant (2026-07-08T05:06:36.312Z)

Let me read the compile-job DB methods and existing tests to mirror the patterns.

31. user (2026-07-08T05:06:36.588Z)

No matches found

32. user (2026-07-08T05:06:39.298Z)

```
=== compile-related tests ===
test_api_cloud_dispatch.py
test_async_jobs.py
test_compiler.py
test_compute_backend.py
test_job_worker_loop.py
test_job_worker_module_seam.py
test_worker.py
=== cov floor ===
.github/workflows/ci.yml:      run: python run_all_tests.py --cov=backend --cov=training
--cov-report=term --cov-fail-under=70
.github/workflows/nightly.yml:  run: python run_all_tests.py --cov=backend --cov=training
--cov-report=term --cov-fail-under=70
```

33. user (2026-07-08T05:06:43.927Z)

```
295:      # re-claims it via `claim_due_job` when due ? with NO master node: the table IS
462:      # v9: the jobs table now holds TWO kinds of work ? `kind`='process' (the original,
NULL ?
463:      # process for pre-v9 rows) and `kind`='compile' (the async reel stitch, #329;
/api/compile
465:      # is the kind-specific JSON payload a compile job carries
(segment_ids/output_filename/locale)
467:      self._add_column_idempotent(cursor, "ALTER TABLE jobs ADD COLUMN kind TEXT")
```

34. user (2026-07-08T05:06:47.170Z)

Found 2 files
backend\infrastructure\database.py
backend\infrastructure\database_jobs.py

35. user (2026-07-08T05:06:49.810Z)

```
1      import json
2      import logging
3      from typing import Any, Dict, List, Optional, Tuple
4
5      from backend import billing # pure cost math (no backend imports -> no cycle)
6
7      logger = logging.getLogger(__name__)
8
9
```

```

10     class DatabaseJobsMixin:
11         def _connect(self): # pragma: no cover - implemented by Database facade
12             raise NotImplementedError
13
14         def _execute_write( # pragma: no cover - implemented by Database facade
15             self, sql: str, params: Tuple[Any, ...], error_msg: str
16         ) -> bool:
17             raise NotImplementedError
18
19         # --- Async job store (DEFERRED_HARDENING_PLAN #16) -----
20         # Writers follow the repo convention (swallow + logger.error, return bool); the
21         # worker logs loudly on a failed transition so a job can't silently wedge.
22
23         def create_job(
24             self,
25             job_id: str,
26             video_path: str,
27             segmenter: Optional[str] = None,
28             player_pool: Optional[List[str]] = None,
29             max_frames: Optional[int] = None,
30             policy: Optional[Dict[str, Any]] = None,
31             owner_id: Optional[str] = None,
32             locale: Optional[str] = None,
33             compute: Optional[str] = None,
34             force: bool = False,
35         ) -> bool:
36             """Insert a new job in state 'queued'. ``policy`` is an optional JSON
ExecutionPolicy
37             (EXECUTION_POLICY.md) the worker honours; None ? the worker's default policy.
``owner_id``
38             (M9, D9) is the verified tenant from edge auth ? NULL for the local single-user
(default).
39             ``locale`` (v7) is the UI language the submission came from ? NULL when
unrecorded.
40             ``compute`` (v8) is the SPA compute-picker choice ("local"/"cloud") the worker
reconstructs +
41             honours; NULL ? the server default (byte-identical to pre-picker behaviour).
``force`` opts
42             into deliberate reprocessing; false leaves params NULL for old-row
compatibility."""
43             params = json.dumps({"force": True}) if force else None
44             return self._execute_write(
45                 "INSERT INTO jobs (id, video_path, segmenter, player_pool, max_frames,
status, policy, "
46                 "owner_id, locale, compute, params) VALUES (?, ?, ?, ?, ?, 'queued', ?, ?,
?, ?, ?)",
47                 (
48                     job_id,
49                     video_path,
50                     segmenter,
51                     json.dumps(player_pool) if player_pool is not None else None,
52                     max_frames,
53                     json.dumps(policy) if policy is not None else None,
54                     owner_id,
55                     locale,
56                     compute,
57                     params,
58                 ),
59                 f"Failed to create job {job_id}",
60             )
61
62         def create_compile_job(
63             self,
64             job_id: str,
65             video_id: str,

```



```

66         video_path: str,
67         segment_ids: List[int],
68         output_filename: str,
69         locale: Optional[str] = None,
70         owner_id: Optional[str] = None,
71     ) -> bool:
72         """Insert an async REEL-COMPILE job (#329) in state 'queued', `kind`='compile'.
The stitch
73         params (segment_ids / output_filename / locale) ride the JSON `params` column
since they don't
74         map onto the process columns; `video_path` is the target video's source ref (the
table's
75         NOT NULL column) and `video_id` is set up-front (so /api/jobs surfaces it
immediately). The
76         worker dispatches on `kind` and runs the ffmpeg stitch off the request
thread."""
77         params = json.dumps(
78             {
79                 "segment_ids": list(segment_ids),
80                 "output_filename": output_filename,
81                 "locale": locale,
82             }
83         )
84         return self._execute_write(
85             "INSERT INTO jobs (id, video_path, video_id, status, owner_id, locale, kind,
params) "
86             "VALUES (?, ?, ?, 'queued', ?, ?, 'compile', ?)",
87             (job_id, video_path, video_id, owner_id, locale, params),
88             f"Failed to create compile job {job_id}",
89         )
90
91     def count_jobs_by_locale(self) -> Dict[str, int]:
92         """PMF analytics: how many jobs were submitted per UI locale. Unrecorded locales
(NULL ?
93         CLI / pre-v7 / older clients) bucket under '(unset)'. Returns {locale: count},
busiest first."""
94         with self._connect() as conn:
95             rows = (
96                 conn.cursor()
97                 .execute(
98                     "SELECT COALESCE(locale, '(unset)') AS loc, COUNT(*) AS n "
99                     "FROM jobs GROUP BY loc ORDER BY n DESC, loc"
100                 )
101                 .fetchall()
102             )
103             return {str(r[0]): int(r[1]) for r in rows}
104
105     def mark_job_running(self, job_id: str) -> bool:
106         return self._execute_write(
107             "UPDATE jobs SET status='running', started_at=CURRENT_TIMESTAMP, "
108             "progress='starting' WHERE id = ?",
109             (job_id,),
110             f"Failed to mark job {job_id} running",
111         )
112
113     def set_job_progress(self, job_id: str, progress: str) -> bool:
114         return self._execute_write(
115             "UPDATE jobs SET progress=? WHERE id = ?",
116             (progress, job_id),
117             f"Failed to set progress for job {job_id}",
118         )
119
120     def set_job_video_id(self, job_id: str, video_id: str) -> bool:
121         return self._execute_write(
122             "UPDATE jobs SET video_id=? WHERE id = ?",

```

```

123         (video_id, job_id),
124         f"Failed to set video_id for job {job_id}",
125     )
126
127 def mark_job_done(self, job_id: str, result: Dict[str, Any]) -> bool:
128     """Job EXECUTED to completion. The pipeline's own verdict (success/failed
129     validation) is inside `result` ? job status 'done' means 'the worker
130     finished'."""
131     return self._execute_write(
132         "UPDATE jobs SET status='done', progress='finished', result=?, "
133         "finished_at=CURRENT_TIMESTAMP WHERE id = ?",
134         (json.dumps(result), job_id),
135         f"Failed to mark job {job_id} done",
136     )
137
138 def mark_job_failed(self, job_id: str, error: str) -> bool:
139     """Worker crashed / raised ? infrastructure failure, not a pipeline verdict."""
140     return self._execute_write(
141         "UPDATE jobs SET status='failed', error=?, "
142         "finished_at=CURRENT_TIMESTAMP WHERE id = ?",
143         (error, job_id),
144         f"Failed to mark job {job_id} failed",
145     )
146
147 def get_job(self, job_id: str) -> Optional[Dict[str, Any]]:
148     """Fetch one job; `result`/`player_pool` JSON deserialized."""
149     with self._connect() as conn:
150         row = (
151             conn.cursor()
152             .execute("SELECT * FROM jobs WHERE id = ?", (job_id,))
153             .fetchone()
154         )
155         if not row:
156             return None
157         d = dict(row)
158         d["result"] = json.loads(d["result"]) if d["result"] else None
159         d["player_pool"] = (
160             json.loads(d["player_pool"]) if d["player_pool"] else None
161         )
162         d["policy"] = json.loads(d["policy"]) if d.get("policy") else None
163         return d
164
165 def list_jobs(
166     self,
167     owner_id: Optional[str] = None,
168     statuses: Optional[List[str]] = None,
169     limit: int = 50,
170 ) -> List[Dict[str, Any]]:
171     """List a tenant's jobs, newest first ? the /api/jobs processing-queue view so a
172     client can
173     discover in-flight work it didn't submit (a new tab/device, cleared storage).
174     ``owner_id``
175     NULL ? the local single-user's NULL-owner rows; a value ? ONLY that tenant's
176     rows (M9, D9
177     isolation ? never leak another tenant's jobs). ``statuses`` filters execution
178     state (e.g.
179     ['queued', 'running', 'waiting'] for the active queue); None ? any state. Light
180     projection: the
181     heavy ``result`` JSON is reduced to ``result_status`` (the pipeline verdict) so
182     the queue stays
183     cheap to poll."""
184     owner_clause = "owner_id IS NULL" if owner_id is None else "owner_id = ?"
185     params: List[Any] = [] if owner_id is None else [owner_id]
186     clauses = [owner_clause]
187     if statuses:

```

```

181         clauses.append("status IN (%s)" % ", ".join("?" for _ in statuses))
182         params.extend(statuses)
183     sql = (
184         "SELECT id, kind, status, progress, video_id, video_path, compute, error,
result, "
185         "created_at, started_at, finished_at FROM jobs WHERE "
186         + " AND ".join(clauses)
187         + " ORDER BY created_at DESC, id DESC LIMIT ?"
188     )
189     params.append(int(limit))
190     with self._connect() as conn:
191         rows = conn.cursor().execute(sql, tuple(params)).fetchall()
192     out: List[Dict[str, Any]] = []
193     for row in rows:
194         d = dict(row)
195         res = json.loads(d["result"]) if d.get("result") else None
196         d["result_status"] = (res or {}).get("status") if res else None
197         d.pop("result", None)
198         out.append(d)
199     return out
200
201     def find_active_process_job(
202         self, owner_id: Optional[str], video_path: str
203     ) -> Optional[Dict[str, Any]]:
204         """Dedup key for /api/process idempotency: the {id, status} of an already-ACTIVE
205         (queued|running|waiting) PROCESS job for the same (owner_id, video_path), else
None ? so a
206         re-submit of the same video (a tester who can't see their running job) returns
that job
207         instead of dispatching a second, PAID GPU run. Compile jobs (kind='compile') are
excluded;
208         the owner scoping mirrors ``list_jobs`` (NULL ? local single-user)."""
209         owner_clause = "owner_id IS NULL" if owner_id is None else "owner_id = ?"
210         params: Tuple[Any, ...] = (
211             (video_path,) if owner_id is None else (owner_id, video_path)
212         )
213         with self._connect() as conn:
214             row = (
215                 conn.cursor()
216                 .execute(
217                     "SELECT id, status FROM jobs WHERE "
218                     + owner_clause
219                     + " AND video_path = ? AND status IN ('queued', 'running',
'waiting') "
220                     + "AND (kind IS NULL OR kind != 'compile') "
221                     + "ORDER BY created_at DESC LIMIT 1",
222                     params,
223                 )
224                 .fetchone()
225             )
226         return {"id": str(row["id"]), "status": str(row["status"])} if row else None
227
228     def claim_or_create_process_job(
229         self,
230         job_id: str,
231         video_path: str,
232         segmenter: Optional[str] = None,
233         player_pool: Optional[List[str]] = None,
234         max_frames: Optional[int] = None,
235         policy: Optional[Dict[str, Any]] = None,
236         owner_id: Optional[str] = None,
237         locale: Optional[str] = None,
238         compute: Optional[str] = None,
239         force: bool = False,
240         video_id: str = "",

```

```

241         ) -> Optional[Dict[str, Any]]:
242             """ATOMIC idempotent submit for a PROCESS job (fixes the read-before-insert race
? PR #481
243             review): in ONE write-locked statement, insert ``job_id`` as 'queued' ONLY IF no
active
244             (queued|running|waiting) process job already exists for (owner_id, video_path).
Two concurrent
245             identical submits (a double-click / retry) therefore can't both insert ? SQLite
serializes
246             writers, and the second's ``INSERT ? WHERE NOT EXISTS`` sees the first's row
(same
247             compare-and-set idiom as ``claim_due_job``, no read-then-write gap). ``force``
skips the guard
248             (deliberate reprocess ? always inserts). Returns
``{"job_id","status","created"}`` ? ``created``
249             False means an active job already existed and the caller must NOT dispatch a
second run ? or
250             ``None`` on failure.
251
252             ``video_id`` is the submit-time MD5 when the caller could resolve one cheaply
(#484 ?
253             gs:// object metadata); "" ? unknown-yet (the drain fills it after hashing, as
before).
254             Persisting it at CREATE is what lets a restarted control plane re-associate a
running
255             remote job with its bucket outputs (CONTROL_PLANE_DURABLE_REBUILD $4.2)."""
256             params = json.dumps({"force": True}) if force else None
257             vals = (
258                 job_id,
259                 video_path,
260                 video_id,
261                 segmenter,
262                 json.dumps(player_pool) if player_pool is not None else None,
263                 max_frames,
264                 json.dumps(policy) if policy is not None else None,
265                 owner_id,
266                 locale,
267                 compute,
268                 params,
269             )
270             owner_clause = "owner_id IS NULL" if owner_id is None else "owner_id = ?"
271             guard = () if owner_id is None else (owner_id,)
272             cols = (
273                 "id, video_path, video_id, segmenter, player_pool, max_frames, status,
policy, "
274                 "owner_id, locale, compute, params"
275             )
276             try:
277                 with self._connect() as conn:
278                     cur = conn.cursor()
279                     if force:
280                         cur.execute(
281                             f"INSERT INTO jobs ({cols}) VALUES (?, ?, ?, ?, ?, ?, 'queued',
?, ?, ?, ?, ?)",
282                             vals,
283                         )
284                     conn.commit()
285                     return {"job_id": job_id, "status": "queued", "created": True}
286                     # Conditional insert ? one atomic statement: only if no active process
job exists for
287                     # this (owner, input). SQLite serializes writers, so a racing second
submit sees the row.
288                     cur.execute(
289                         f"INSERT INTO jobs ({cols}) "
290                         "SELECT ?, ?, ?, ?, ?, ?, 'queued', ?, ?, ?, ? WHERE NOT EXISTS

```

```

(
291         "SELECT 1 FROM jobs WHERE " + owner_clause + " AND video_path = ? "
292         "AND status IN ('queued', 'running', 'waiting') "
293         "AND (kind IS NULL OR kind != 'compile'))",
294         vals + guard + (video_path,),
295     )
296     conn.commit()
297     if cur.rowcount == 1:
298         return {"job_id": job_id, "status": "queued", "created": True}
299     # Lost the race (or a pre-existing active job): return the winner, no
second dispatch.
300     row = (
301         cur.execute(
302             "SELECT id, status FROM jobs WHERE " + owner_clause + " AND
video_path = ? "
303             "AND status IN ('queued', 'running', 'waiting') "
304             "AND (kind IS NULL OR kind != 'compile') ORDER BY created_at
DESC LIMIT 1",
305             guard + (video_path,),
306         ).fetchone()
307     )
308     if row:
309         return {"job_id": str(row["id"]), "status": str(row["status"]),
"created": False}
310     return None
311 except Exception as e:
312     logger.error(f"Failed to claim-or-create process job {job_id}: {e}")
313     return None
314
315 # --- Per-job cost ledger (COMPUTE_DECOUPLED_SERVING M8, D8)
-----
316 def get_pricebook_rates(
317     self, version: str, gpu_type: Optional[str] = None
318 ) -> Dict[str, float]:
319     """Resolve a flat ``{resource: unit_price_usd}`` for a pricebook ``version``.
The per-second
320 GPU rate is selected by ``gpu_type``; non-GPU resources (wall/egress) carry
``gpu_type=''``.
321 An unknown version ? ``{}`` (so every cost computes to 0.0)."""
322 gt = gpu_type or ""
323 rates: Dict[str, float] = {}
324 with self._connect() as conn:
325     rows = (
326         conn.cursor()
327         .execute(
328             "SELECT resource, gpu_type, unit_price_usd FROM pricebook WHERE
version = ?",
329             (version,),
330         )
331         .fetchall()
332     )
333     for r in rows:
334         res, row_gt, price = (
335             r["resource"],
336             (r["gpu_type"] or ""),
337             float(r["unit_price_usd"]),
338         )
339         if res == billing.GPU_SECONDS:
340             if row_gt == gt: # the GPU rate for THIS gpu_type
341                 rates[res] = price
342             else:
343                 rates[res] = price # wall/egress: gpu_type-independent
344     return rates
345
346 def record_job_cost(

```

```

347         self,
348         job_id: str,
349         *,
350         usage: Dict[str, float],
351         pricebook_version: str,
352         gpu_type: Optional[str] = None,
353         compute_backend: Optional[str] = None,
354         owner_id: Optional[str] = None,
355         margin: float = 0.0,
356     ) -> Optional[Dict[str, Any]]:
357         """Compute the job's cost from ``usage`` × the pricebook + persist every cost
column.
358         Returns the stored cost dict, or ``None`` on write failure. The caller gates
this on
359         ``billing.enabled`` (default-OFF) ? nothing records until billing is turned
on."""
360         rates = self.get_pricebook_rates(pricebook_version, gpu_type)
361         # SILENT-UNDER-BILL GUARD: a GPU-seconds AMOUNT with no RATE for this gpu_type
would bill at
362         # $0 indistinguishably from a free run (a new SKU / typo / gpu_type=None).
Surface it loudly +
363         # in the breakdown so the gap is auditable, not invisible. (A missing amount =
$0 IS safe.)
364         gpu_amt = float(usage.get(billing.GPU_SECONDS, 0.0) or 0.0)
365         unpriced_gpu = gpu_amt > 0 and billing.GPU_SECONDS not in rates
366         if unpriced_gpu:
367             logger.warning(
368                 "[BILLING] no pricebook rate for gpu_type=%r in version=%r ? %.1f GPU-
seconds billed "
369                 "at $0 (UNPRICED; recorded in cost_breakdown_json). Add a rate or fix
gpu_type.",
370                 gpu_type,
371                 pricebook_version,
372                 gpu_amt,
373             )
374         cost_usd, breakdown = billing.compute_cost_usd(usage, rates)
375         if unpriced_gpu:
376             breakdown["unpriced_gpu_seconds"] = gpu_amt
377         if margin:
378             marked = billing.apply_margin(cost_usd, margin)
379             breakdown["margin"] = round(
380                 marked - cost_usd, 6
381             ) # keep ?(breakdown) == cost_usd
382             cost_usd = marked
383         bytes_out = int(float(usage.get(billing.EGRESS_GB, 0.0) or 0.0) * 1e9)
384         ok = self._execute_write(
385             "UPDATE jobs SET owner_id=?, compute_backend=?, gpu_type=?,
gpu_active_seconds=?, "
386             "wall_seconds=?, bytes_out=?, gemini_cost_usd=?, pricebook_version=?,
cost_usd=?, "
387             "cost_breakdown_json=? WHERE id = ?",
388             (
389                 owner_id,
390                 compute_backend,
391                 gpu_type,
392                 usage.get(billing.GPU_SECONDS),
393                 usage.get(billing.WALL_SECONDS),
394                 bytes_out,
395                 usage.get(billing.GEMINI_USD),
396                 pricebook_version,
397                 cost_usd,
398                 json.dumps(breakdown),
399                 job_id,
400             ),
401         f"Failed to record cost for job {job_id}",

```

```

402         )
403         if not ok:
404             return None
405         return {
406             "cost_usd": cost_usd,
407             "cost_breakdown": breakdown,
408             "pricebook_version": pricebook_version,
409         }
410
411     def get_job_cost(self, job_id: str) -> Optional[Dict[str, Any]]:
412         """The job's cost row (``None`` if the job doesn't exist; ``cost_usd`` is
413         ``None`` until
414         recorded). ``cost_breakdown`` is the deserialized JSON breakdown."""
415         with self._connect() as conn:
416             row = (
417                 conn.cursor()
418                 .execute(
419                     "SELECT id, video_id, owner_id, compute_backend, gpu_type,
420                     gpu_active_seconds, "
421                     "wall_seconds, bytes_out, gemini_cost_usd, pricebook_version,
422                     cost_usd, "
423                     "cost_breakdown_json FROM jobs WHERE id = ?",
424                     (job_id,),
425                 )
426                 .fetchone()
427             )
428         if not row:
429             return None
430         d = dict(row)
431         raw = d.pop("cost_breakdown_json", None)
432         d["cost_breakdown"] = json.loads(raw) if raw else {}
433         return d
434
435     def get_video_cost(self, video_id: str) -> Dict[str, Any]:
436         """Aggregate cost across a video's jobs: total ``cost_usd`` + the per-job rows.
437         A video is
438         1:1 with an infer job, but a re-ingest can produce several ? sum them."""
439         with self._connect() as conn:
440             rows = (
441                 conn.cursor()
442                 .execute(
443                     "SELECT id, cost_usd, pricebook_version, cost_breakdown_json FROM
444                     jobs "
445                     "WHERE video_id = ?",
446                     (video_id,),
447                 )
448                 .fetchall()
449             )
450         jobs: List[Dict[str, Any]] = []
451         total = 0.0
452         for r in rows:
453             cost = r["cost_usd"]
454             if cost is not None:
455                 total += float(cost)
456             jobs.append(
457                 {
458                     "job_id": r["id"],
459                     "cost_usd": cost,
460                     "pricebook_version": r["pricebook_version"],
461                     "cost_breakdown": json.loads(r["cost_breakdown_json"])
462                     if r["cost_breakdown_json"]
463                     else {},
464                 }
465             )
466         return {"video_id": video_id, "total_cost_usd": round(total, 6), "jobs": jobs}

```

```

462
463         # --- Self-scheduling execution policy (EXECUTION_POLICY.md P2)
-----
464     def set_job_policy(self, job_id: str, policy: Optional[Dict[str, Any]]) -> bool:
465         """Attach/replace a job's JSON ExecutionPolicy."""
466         return self._execute_write(
467             "UPDATE jobs SET policy=? WHERE id = ?",
468             (json.dumps(policy) if policy is not None else None, job_id),
469             f"Failed to set policy for job {job_id}",
470         )
471
472     def set_job_waiting(
473         self, job_id: str, delay_sec: float, progress: Optional[str] = None
474     ) -> bool:
475         """Self-schedule: park a job in 'waiting' until ``delay_sec`` from now,
releasing the
476         worker. `claim_due_job` re-picks it when due ? cooperative, no master.
``next_poll_at`` is
477         computed in SQLite's own UTC format so it compares directly to
CURRENT_TIMESTAMP."""
478         return self._execute_write(
479             "UPDATE jobs SET status='waiting', "
480             "next_poll_at=datetime('now', ?), "
481             "progress=COALESCE(?, progress) WHERE id = ?",
482             (f"+{max(0, int(delay_sec))} seconds", progress, job_id),
483             f"Failed to set job {job_id} waiting",
484         )
485
486     def seconds_until_next_due(self) -> Optional[float]:
487         """Seconds until the soonest parked ('waiting') job becomes due, so the worker
can
488         schedule a wake-up (EXECUTION_POLICY.md P3 self-scheduling). Returns 0.0 if one
is
489         already due, or None if nothing is waiting. A NULL ``next_poll_at`` ? due now
(0.0)."""
490         try:
491             with self._connect() as conn:
492                 row = (
493                     conn.cursor()
494                     .execute(
495                         "SELECT MIN(CASE WHEN next_poll_at IS NULL THEN 0 ELSE "
496                         "    (julianday(next_poll_at) - julianday('now')) * 86400.0 END)
AS secs "
497                         "FROM jobs WHERE status='waiting'"
498                     )
499                     .fetchone()
500                 )
501                 if row is None or row["secs"] is None:
502                     return None
503                 return max(0.0, float(row["secs"]))
504         except Exception as e:
505             logger.error(f"Failed to compute next-due delay: {e}")
506             return None
507
508     def claim_due_job(self) -> Optional[Dict[str, Any]]:
509         """Atomically claim ONE runnable job ? 'queued', or 'waiting' whose
``next_poll_at`` is due ?
510         flipping it to 'running'. Compare-and-set on (id, status) so concurrent workers
never
511         double-claim (the DB row-claim is the only coordination needed; no master).
Earliest-due
512         first. Returns the claimed job dict, or None ONLY when a fresh SELECT finds
nothing
513         runnable ? never for a lost race.
514

```



```

515         A lost CAS RETRIES instead of returning: under the parallel drain (#466) several
workers
516         can SELECT the same head-of-queue row, and if the losers reported None the
caller
517         (`_drain_due_jobs`) would read it as "queue empty" and stop with runnable jobs
still
518         queued ? stranding dispatch capacity until some later submit/wake. The retry
loop is
519         deliberately UNBOUNDED: every lost race means another worker successfully
claimed that
520         row, so the runnable set shrank and global progress is guaranteed (lock-free
progress);
521         a retry cap would just reintroduce the false-empty. One connection serves all
attempts ?
522         the SELECT is an autocommit read and each CAS commits before the next iteration,
so no
523         write lock is ever held across a retry (busy_timeout covers momentary
contention)."""
524         try:
525             with self._connect() as conn:
526                 cur = conn.cursor()
527                 while True:
528                     row = cur.execute(
529                         "SELECT id, status FROM jobs WHERE status='queued' "
530                         "OR (status='waiting' AND (next_poll_at IS NULL OR next_poll_at
<= CURRENT_TIMESTAMP)) "
531                         "ORDER BY COALESCE(next_poll_at, created_at) ASC LIMIT 1"
532                     ).fetchone()
533                     if not row:
534                         return None # nothing runnable ? a TRUE empty, not a lost race
535                     cur.execute(
536                         "UPDATE jobs SET status='running', started_at=CURRENT_TIMESTAMP
"
537                         "WHERE id=? AND status=?",
538                         (row["id"], row["status"]),
539                     )
540                     conn.commit()
541                     if cur.rowcount == 1:
542                         break # claimed it
543                     # Lost the race: the winner flipped this row; re-SELECT for the next
one.
544                 return self.get_job(row["id"])
545         except Exception as e:
546             logger.error(f"Failed to claim a due job: {e}")
547         return None
548
549     def fail_stale_jobs(self, reason: str = "interrupted by server restart") -> int:
550         """Crash recovery: any job still queued/running from a previous process is dead
551         (the executor is in-process). Mark failed so clients see a terminal state, not a
552         hang. Returns the number of jobs swept.
553
554         NOTE: ``recover_stale_jobs`` is the preferred startup call ? it RE-QUEUES stale
555         ``compile`` jobs (idempotent CPU stitch, safe to retry) instead of failing them.
This
556         blanket-fail method stays for callers/tests that want the old behaviour."""
557         try:
558             with self._connect() as conn:
559                 cur = conn.cursor()
560                 cur.execute(
561                     "UPDATE jobs SET status='failed', error=?,
finished_at=CURRENT_TIMESTAMP "
562                     "WHERE status IN ('queued', 'running')",
563                     (reason,),
564                 )
565                 conn.commit()

```

```

566         return cur.rowcount
567     except Exception as e:
568         logger.error(f"Failed to sweep stale jobs: {e}")
569         return 0
570
571     def recover_stale_jobs(
572         self, reason: str = "interrupted by server restart"
573     ) -> Tuple[int, int]:
574         """Crash recovery, kind-aware (#329 serving reliability). Called once at
startup.
575
576         A ``compile`` job is a **CPU-only, idempotent** ffmpeg re-stitch (re-resolving
the same
577         segments yields the same reel), so a compile left queued/running by a killed
instance is
578         **re-queued** ? the startup drain re-runs it and the client keeps polling until
it gets
579         the reel, instead of seeing a spurious "failed". A ``process`` job is an
expensive,
580         non-idempotent GPU detection pass; auto-retrying could double-spend, so it is
**failed**
581         (terminal, as before) and the client re-submits deliberately ? EXCEPT a RUNNING
process
582         job whose ``video_id`` is already known (CP-rebuild P3): its remote worker may
still be
583         computing, and ``outputs/<md5>/report.json`` is a durable done-signal the
restarted
584         control plane can re-associate by md5. Those rows are LEFT RUNNING for
585         ``job_worker.resume_interrupted_remote_jobs`` to complete or fail honestly ?
never
586         re-dispatched (no double GPU spend). This is exactly the 2026-07-05 loss shape:
the
587         worker finished 45 min after the restart and its 29 rallies had no listener.
588
589         ``force=true`` rows are NOT resumable (#494 review): force deliberately bypasses
the
590         completed-work cache, and without a persisted per-dispatch identity the bucket
report
591         cannot be proven to belong to the forced run rather than an older completed one
? so a
592         resumed forced job could be marked done from exactly the stale result the user
asked
593         to bypass. Forced rows keep today's honest terminal failure on restart.
594
595         ``compute='local'`` rows are NOT resumable either (#494 review round 2): an
explicit
596         local override forces the in-process segmenter ? no remote worker exists, and a
bucket
597         report under the same md5 would be an OLDER run's output, not this job's.
``compute``
598         NULL (server default) stays in the carve-out at the SQL level; the resume pass ?
which
599         has the config ? fails it unless the server default target is actually
cloud_run.
600
601         Returns ``(requeued_compile, failed_process)``.
602         try:
603             with self._connect() as conn:
604                 cur = conn.cursor()
605                 # Re-queue interrupted COMPILE jobs (idempotent ? safe to retry
automatically).
606                 cur.execute(
607                     "UPDATE jobs SET status='queued', started_at=NULL,
next_poll_at=NULL, "
608                     "error=NULL WHERE status IN ('queued', 'running') AND

```

```

kind='compile'"
609         )
610         requeued = cur.rowcount
611         # Fail interrupted PROCESS jobs (GPU, not safe to auto-retry) ? except
the
612         # md5-keyed RUNNING ones, which the resume pass owns (see docstring).
613         cur.execute(
614             "UPDATE jobs SET status='failed', error=?,
finished_at=CURRENT_TIMESTAMP "
615             "WHERE status IN ('queued', 'running') AND (kind IS NULL OR kind !=
'compile') "
616             "AND NOT (status = 'running' AND video_id IS NOT NULL AND video_id
!= ' ' "
617             "AND COALESCE(json_extract(params, '$.force'), 0) != 1 "
618             "AND (compute IS NULL OR LOWER(TRIM(compute)) != 'local'))",
619             (reason, ),
620         )
621         failed = cur.rowcount
622         conn.commit()
623         return requeued, failed
624     except Exception as e:
625         logger.error(f"Failed to recover stale jobs: {e}")
626         return 0, 0
627
628     def interrupted_remote_process_jobs(self) -> List[Dict[str, Any]]:
629         """The RUNNING md5-keyed process jobs a restart left behind ? the rows
630         ``recover_stale_jobs`` deliberately did NOT fail (CP-rebuild P3). Called once at
631         startup, BEFORE the drain runs, so every 'running' process row here is orphaned
by
632         definition (the executor is in-process; no live worker thread can exist yet)."""
633         try:
634             with self._connect() as conn:
635                 rows = conn.execute(
636                     "SELECT id, video_id, video_path, owner_id, compute FROM jobs "
637                     "WHERE status = 'running' AND (kind IS NULL OR kind != 'compile') "
638                     "AND video_id IS NOT NULL AND video_id != ' ' "
639                     "AND COALESCE(json_extract(params, '$.force'), 0) != 1 "
640                     "AND (compute IS NULL OR LOWER(TRIM(compute)) != 'local'))"
641                 ).fetchall()
642                 return [dict(r) for r in rows]
643         except Exception as e:
644             logger.error(f"Failed to list interrupted remote process jobs: {e}")
645             return []
646
647

```

36. assistant (2026-07-08T05:11:57.940Z)

I now understand the architecture well. Let me read the existing tests to mirror their patterns before implementing.

37. user (2026-07-08T05:11:58.949Z)

```

1         """API cloud-serving dispatch (COMPUTE_DECOUPLED_SERVING M6 ? the API path).
2
3         Proves the /api/process job path can offload DETECT to a Cloud Run GPU Job and ingest
the rally
4         intervals back into the API DB, so /api/query + /api/compile (local stitch) are
unchanged ? all
5         GPU-free + cloud-free via a FAKE CloudRunJobClient that writes report.json +
intervals.csv + the
6         done-marker into a local "bucket". Cardinal invariant: DEFAULT-OFF ? only
compute_target=cloud_run
7         dispatches; every other target runs the in-process segmenter byte-identically.
8         """

```

```

9
10 import csv
11 import json
12 import os
13
14 import pytest
15
16 pytest.importorskip("httpx")
17 from fastapi.testclient import TestClient
18
19 import backend.compute as compute_pkg
20 import backend.main as m
21 from backend.compute.cloud_run import CloudRunJobClient
22 from backend.config.models import AppConfig
23 from backend.infrastructure.database import Database
24 from backend.infrastructure.execution_policy import ExecutionPolicy
25 from backend.pipeline.validators.base import ValidationResult
26
27 _FAST = ExecutionPolicy(poll_every_n_sec=0.0, op_timeout_sec=5.0, max_wait_sec=5.0)
28
29
30 def _pass():
31     return [ValidationResult(validator_name="x", passed=True, message="ok", details={})]
32
33
34 class _InlineExecutor:
35     def submit(self, fn, *a, **k):
36         fn(*a, **k)
37
38
39 class _FakeCloudClient(CloudRunJobClient):
40     """Simulates the Cloud Run worker: on dispatch, write report.json + intervals.csv +
the done-marker
41     into the local bucket exactly where the real worker would, so the bucket-poll +
ingest resolve."""
42
43     def __init__(self, *, video_id="vidCloud", returncode=0, rows=None):
44         self.calls: list[list[str]] = []
45         self.video_id, self.returncode = video_id, returncode
46         self.rows = (
47             rows
48             if rows is not None
49             else [(2.74, 6.91, 5, "winner"), (8.74, 11.38, 3, "unforced_error")]
50         )
51
52     def run_job(self, job_name, argv, *, region, project=None):
53         self.calls.append(list(argv))
54         out_uri = argv[argv.index("--out-uri") + 1].rstrip("/")
55         done_uri = argv[argv.index("--done-uri") + 1]
56         if self.returncode == 0:
57             outputs = os.path.join(out_uri, "outputs", self.video_id)
58             os.makedirs(outputs, exist_ok=True)
59             with open(os.path.join(outputs, "report.json"), "w", encoding="utf-8") as f:
60                 json.dump(
61                     {
62                         "video_id": self.video_id,
63                         "segment_count": len(self.rows),
64                         "status": "success",
65                     },
66                     f,
67                 )
68             with open(
69                 os.path.join(outputs, "intervals.csv"),
70                 "w",
71                 newline="",

```

```

72         encoding="utf-8",
73     ) as f:
74         w = csv.writer(f)
75         w.writerow(["start", "end", "shot_count", "ending_reason"])
76         for s, e, sc, er in self.rows:
77             w.writerow([f"{s:.3f}", f"{e:.3f}", sc, er])
78     os.makedirs(os.path.dirname(done_uri), exist_ok=True)
79     with open(done_uri, "w", encoding="utf-8") as f:
80         json.dump(
81             {
82                 "video_id": self.video_id,
83                 "returncode": self.returncode,
84                 "segment_count": len(self.rows),
85                 "status": "success",
86             },
87             f,
88         )
89     return "exec-1"
90
91     def cancel_execution(self, execution):
92         raise AssertionError(
93             "cancel must never fire on the happy path"
94         ) # pragma: no cover
95
96
97     @pytest.fixture
98     def cloud_env(tmp_path, monkeypatch):
99         """API wired for cloud_run with a LOCAL bucket (the local storage backend makes the
100         bucket read an
101         identity file read ? no GCS). A gs:// input is 'fetched' to a local fake file for
102         hash+validate."""
103         db = Database(str(tmp_path / "t.db")) # noqa: F841
104         monkeypatch.setattr(m, "db_instance", db)
105         monkeypatch.setattr(m, "db_instance", db)
106         bucket = tmp_path / "bucket"
107         bucket.mkdir()
108         cfg = AppConfig.model_validate(
109             { # noqa: F841
110                 "deployment": {"compute_target": "cloud_run"},
111                 "storage": {"backend": "local", "output_root": str(bucket)},
112             }
113         )
114         monkeypatch.setattr(m, "global_config", cfg)
115         monkeypatch.setattr(m, "global_config", cfg)
116         monkeypatch.setattr(m, "_get_executor", lambda: _InlineExecutor())
117         # A real gs:// input would be fetched to hash+validate; return a local fake so that
118         path is offline.
119         fake_local = tmp_path / "fetched.mp4"
120         fake_local.write_bytes(b"FAKEVIDEOBYTES")
121         monkeypatch.setattr(
122             m.storage, "acquire_local_input", lambda path, scfg: str(fake_local)
123         )
124         monkeypatch.setattr(
125             m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {})
126         )
127         return db, bucket, monkeypatch
128
129     def _inject_fake(monkeypatch, fake):
130         """Make backend.compute.get_compute_backend (re-imported inside
131         _maybe_dispatch_remote) build a
132         real CloudRunCompute around the FAKE client + the fast no-wait policy + a fixed
133         token. ``**kw``
134         forwards the per-request ``target_override`` (item 3) so the cloud-picker path
135         resolves too."""

```

```

131     real = compute_pkg.get_compute_backend
132     monkeypatch.setattr(
133         compute_pkg,
134         "get_compute_backend",
135         lambda config, **kw: real(
136             config, run_client=fake, policy=_FAST, token="tok", **kw
137         ),
138     )
139
140
141 def _meta(row):
142     md = row.get("metadata")
143     return json.loads(md) if isinstance(md, str) else (md or {})
144
145
146 def test_cloud_dispatch_ingests_segments(cloud_env):
147     db, _bucket, monkeypatch = cloud_env
148     fake = _FakeCloudClient()
149     _inject_fake(monkeypatch, fake)
150
151     out = m._execute_processing(
152         m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
153     )
154
155     assert out["status"] == "success"
156     assert len(out["play_segments"]) == 2
157     rows = db.query_play_segments(out["video_id"], {})
158     assert len(rows) == 2
159     metas = [_meta(r) for r in rows]
160     assert all(md["source"] == "cloud_run" for md in metas)
161     assert {md.get("shot_count") for md in metas} == {5, 3}
162     assert {md.get("ending_reason") for md in metas} == {"winner", "unforced_error"}
163
164
165 def test_cloud_dispatch_builds_correct_spec(cloud_env):
166     _db, bucket, monkeypatch = cloud_env
167     fake = _FakeCloudClient()
168     _inject_fake(monkeypatch, fake)
169
170     m._execute_processing(
171         m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
172     )
173
174     assert len(fake.calls) == 1 # exactly one Cloud Run Job execution
175     argv = fake.calls[0]
176     assert argv[argv.index("--video-uri") + 1] == "gs://b/clip.mp4"
177     assert argv[argv.index("--out-uri") + 1] == str(bucket)
178     assert argv[argv.index("--segmenter") + 1] == "wasb_hybrid"
179     joined = " ".join(argv)
180     assert (
181         "deployment.compute_target=local_gpu" in joined
182     ) # container runs INLINE (never re-dispatch)
183     assert "indexing.detector_impl=native" in joined
184     assert (
185         "indexing.skip_ai_handoff=false" in joined
186     ) # the engine's AI-handoff stance is forwarded
187
188
189 def test_cloud_dispatch_forwards_skip_ai_handoff_true(cloud_env):
190     """When the engine runs no-AI (skip_ai_handoff=True), the cloud job must too ? else
191     it tries the
192     Gemini handoff (which the cloud job has no key for) and yields 0 rallies. The flag
193     is forwarded."""
194     _db, _bucket, monkeypatch = cloud_env
195     m.global_config.indexing.skip_ai_handoff = True

```

```

194     fake = _FakeCloudClient()
195     _inject_fake(monkeypatch, fake)
196     m._execute_processing(
197         m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
198     )
199     assert "indexing.skip_ai_handoff=true" in " ".join(fake.calls[0])
200
201
202     def test_cloud_dispatch_forwards_wasb_serving_knobs(cloud_env):
203         """The worker Job has no DEPLOYMENT_PROFILE, so it never applies the cloud-serving
204         preset ?
205         the control-plane MUST forward its resolved WASB serving knobs (decode_backend +
206         GPU-feed
207         tuning) or the worker silently reverts to the model defaults (cpu decode, batch 8),
208         making the
209         preset flip (#506/#507) a no-op on the real dispatch path. Regression guard for that
210         wiring."""
211         _db, _bucket, monkeypatch = cloud_env
212         m.global_config.indexing.wasb.decode_backend = "nvdec_resident"
213         m.global_config.indexing.wasb.batch_size = 64
214         m.global_config.indexing.wasb.prefetch_batches = 4
215         fake = _FakeCloudClient()
216         _inject_fake(monkeypatch, fake)
217         m._execute_processing(
218             m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
219         )
220         joined = " ".join(fake.calls[0])
221         assert "indexing.wasb.decode_backend=nvdec_resident" in joined
222         assert "indexing.wasb.batch_size=64" in joined
223         assert "indexing.wasb.prefetch_batches=4" in joined
224
225     def test_cloud_dispatch_omits_unset_wasb_tuning(cloud_env):
226         """decode_backend defaults to cpu (always forwarded ? harmless, keeps worker ==
227         control-plane),
228         but batch_size/prefetch_batches default to None and must NOT be forwarded (None =
229         the worker's
230         parity default; forwarding an empty value would reach argparse as a bad token)."""
231         _db, _bucket, monkeypatch = cloud_env
232         # defaults: decode_backend="cpu", batch_size=None, prefetch_batches=None
233         fake = _FakeCloudClient()
234         _inject_fake(monkeypatch, fake)
235         m._execute_processing(
236             m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
237         )
238         joined = " ".join(fake.calls[0])
239         assert "indexing.wasb.decode_backend=cpu" in joined
240         assert "indexing.wasb.batch_size" not in joined
241         assert "indexing.wasb.prefetch_batches" not in joined
242
243     def test_cloud_dispatch_skips_worker_revalidation(cloud_env):
244         """The control-plane is the AUTHORITATIVE validation gate: it validated + only
245         dispatches on PASS
246         (registry.run_all_with_context is stubbed to pass in cloud_env). Re-validating on
247         the worker is
248         redundant, decodes the video a second time, and ? because the control-plane's
249         validation config
250         isn't forwarded ? can REJECT a clip the gate already passed (observed 2026-07-07:
251         passed at
252         validation.min_blur_threshold=8.0, then worker-rejected at its default 20.0). So the
253         dispatch MUST
254         carry --skip-validation, telling the worker to trust the gate. Regression guard for
255         that wiring."""
256         _db, _bucket, monkeypatch = cloud_env

```

```

247     fake = _FakeCloudClient()
248     _inject_fake(monkeypatch, fake)
249     m._execute_processing(
250         m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
251     )
252     assert "--skip-validation" in fake.calls[0]
253
254
255     def test_cloud_failure_marks_failed_no_rows(cloud_env):
256         db, _bucket, monkeypatch = cloud_env
257         # pre-seed a prior good row ? must be preserved on a cloud failure (destroy-AFTER-
confirm).
258         db.add_video("vidLocalPrior", "gs://b/clip.mp4", "processed", [])
259         fake = _FakeCloudClient(returncode=1)
260         _inject_fake(monkeypatch, fake)
261
262         out = m._execute_processing(
263             m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
264         )
265         assert out["status"] == "failed"
266         assert "returncode 1" in out["message"]
267         assert out["play_segments"] == []
268         assert db.query_play_segments(out["video_id"], {}) == []
269         # It DID dispatch + run remotely (then failed rc=1), so the path is honestly
cloud_run + dispatched.
270         assert (
271             out["compute"]["path"] == "cloud_run" and out["compute"]["dispatched"] is True
272         )
273
274
275     def test_cloud_requires_a_bucket_input(cloud_env, tmp_path):
276         _db, _bucket, monkeypatch = cloud_env
277         fake = _FakeCloudClient()
278         _inject_fake(monkeypatch, fake)
279         local_path = str(tmp_path / "local_only.mp4")
280
281         out = m._execute_processing(
282             m.ProcessRequest(video_path=local_path, segmenter="wasb_hybrid")
283         )
284         assert out["status"] == "failed"
285         assert "cloud input" in out["message"]
286         assert fake.calls == [] # never dispatched a local-only path
287
288
289     def test_cloud_stages_local_upload_to_bucket(cloud_env, tmp_path, monkeypatch):
290         """A browser upload lands on THIS box's local disk; under cloud-serving with a CLOUD
input_root it
291         is STAGED to the bucket and the L4 job dispatches from the staged gs:// URI (not
hard-failed) ? #461.
292         The DB video record is repointed at the staged source so a later /api/compile can
re-fetch it."""
293         db, _bucket, _mp = cloud_env
294         m.global_config.storage.input_backend = "gcs"
295         m.global_config.storage.input_root = "gs://in-bucket/videos"
296         puts: list = []
297
298         def _fake_put(src, uri, **kw):
299             puts.append((str(src), uri, kw.get("config")))
300             return uri
301
302         monkeypatch.setattr(m.storage, "put", _fake_put)
303         fake = _FakeCloudClient()
304         _inject_fake(monkeypatch, fake)
305         local_path = str(tmp_path / "upload.mp4")
306

```



```

307         out = m._execute_processing(
308             m.ProcessRequest(video_path=local_path, segmenter="wasb_hybrid")
309         )
310
311         assert out["status"] == "success"
312         # 1) the local upload was staged UP to the input bucket, namespaced by video_id
313         assert len(puts) == 1
314         src, dest, put_cfg = puts[0]
315         assert src == local_path
316         assert dest.startswith("gs://in-bucket/videos/") and dest.endswith("/upload.mp4")
317         # the active storage settings are threaded into put() (so a non-default S3 endpoint
/ GCS project /
318         # etc. survives) ? not a fresh default client. The model_dump carries
input_root/input_backend.
319         assert put_cfg is not None and put_cfg.get("input_root") == "gs://in-bucket/videos"
320         assert put_cfg.get("input_backend") == "gcs"
321         # 2) the L4 job dispatched from the STAGED gs:// URI, not the unreadable local path
322         assert len(fake.calls) == 1
323         argv = fake.calls[0]
324         assert argv[argv.index("--video-uri") + 1] == dest
325         # 3) the DB video record now points at the staged source (so compile re-fetches it)
326         row = db.get_video(out["video_id"])
327         assert row is not None and row["filepath"] == dest
328
329
330     def test_cloud_stages_when_input_root_is_cloud_and_input_backend_default(
331         cloud_env, tmp_path, monkeypatch
332     ):
333         """input_root='gs://?' with input_backend left at its DEFAULT ('local') is a valid
cloud config ?
334         the storage layer classifies an explicit root by its scheme (resolve_input_uri joins
under
335         input_root before ever consulting input_backend), so the staging gate must too (PR
#467 review).
336         Regression: the gate previously required input_backend != 'local' and hard-failed
this shape."""
337         _db, _bucket, _mp = cloud_env
338         m.global_config.storage.input_root = "gs://in-bucket/videos"
339         # input_backend deliberately NOT set ? stays at the model default 'local'
340         assert m.global_config.storage.input_backend == "local"
341         puts: list = []
342
343         def _fake_put(src, uri, **kw):
344             puts.append((str(src), uri))
345             return uri
346
347         monkeypatch.setattr(m.storage, "put", _fake_put)
348         fake = _FakeCloudClient()
349         _inject_fake(monkeypatch, fake)
350         local_path = str(tmp_path / "upload.mp4")
351
352         out = m._execute_processing(
353             m.ProcessRequest(video_path=local_path, segmenter="wasb_hybrid")
354         )
355
356         assert out["status"] == "success" # NOT the 'needs a cloud input' pre-dispatch
failure
357         assert len(puts) == 1
358         src, dest = puts[0]
359         assert src == local_path
360         assert dest.startswith("gs://in-bucket/videos/") and dest.endswith("/upload.mp4")
361         # dispatch continued from the STAGED gs:// URI
362         assert len(fake.calls) == 1
363         argv = fake.calls[0]
364         assert argv[argv.index("--video-uri") + 1] == dest

```

```

365
366
367     def test_cloud_local_upload_without_input_root_still_fails(cloud_env, tmp_path,
monkeypatch):
368         """No cloud input_root configured ? we genuinely can't stage a local upload ? hard-
fail (no
369         dispatch), rather than dispatching a path the Cloud Run job can't read."""
370         _db, _bucket, _mp = cloud_env # fixture leaves input_root unset
(input_backend=local)
371         fake = _FakeCloudClient()
372         _inject_fake(monkeypatch, fake)
373         called = []
374         monkeypatch.setattr(m.storage, "put", lambda *a, **k: called.append(a))
375
376         out = m._execute_processing(
377             m.ProcessRequest(video_path=str(tmp_path / "up.mp4"), segmenter="wasb_hybrid")
378         )
379         assert out["status"] == "failed"
380         assert "cloud input" in out["message"]
381         assert fake.calls == [] and called == [] # neither staged nor dispatched
382
383
384     def test_default_off_runs_in_process(tmp_path, monkeypatch):
385         """compute_target='' ? get_compute_backend returns None ? the in-process segmenter
runs
386         byte-identically (the gate is inert). No Cloud Run dispatch."""
387         db = Database(str(tmp_path / "t.db")) # noqa: F841
388         monkeypatch.setattr(m, "db_instance", db)
389         monkeypatch.setattr(m, "db_instance", db)
390         cfg = AppConfig() # compute_target defaults to "" # noqa: F841
391         monkeypatch.setattr(m, "global_config", cfg)
392         monkeypatch.setattr(m, "global_config", cfg)
393         monkeypatch.setattr(m, "global_config", cfg)
394         monkeypatch.setattr(
395             m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
396         )
397         video = tmp_path / "v.mp4"
398         video.write_bytes(b"hashable-bytes")
399
400         class _InProcSeg:
401             def __init__(self, db_, cfg_):
402                 self.db = db_
403
404             def process_video(self, vid, path, *a, **k):
405                 sid = self.db.add_play_segment(vid, 0.0, 5.0)
406                 return [
407                     {
408                         "id": sid,
409                         "start_time": 0.0,
410                         "end_time": 5.0,
411                         "duration": 5.0,
412                         "metadata": {},
413                     }
414                 ]
415
416         monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _InProcSeg)
417         out = m._execute_processing(
418             m.ProcessRequest(video_path=str(video), segmenter="motion")
419         )
420         assert out["status"] == "success"
421         assert len(out["play_segments"]) == 1
422         assert "cloud_run" not in out["message"]
423
424
425     def test_cloud_ingest_failure_marks_failed_no_leak(cloud_env):

```

```

426         """A rc=0 marker but a MISSING intervals.csv (a partial push) must fail cleanly +
leak no rows
427         (destroy-AFTER-confirm), not 500 or hang."""
428         db, _bucket, monkeypatch = cloud_env
429
430         class _NoCsvClient(_FakeCloudClient):
431             def run_job(self, job_name, argv, *, region, project=None):
432                 self.calls.append(list(argv))
433                 out_uri = argv[argv.index("--out-uri") + 1].rstrip("/")
434                 done_uri = argv[argv.index("--done-uri") + 1]
435                 outputs = os.path.join(out_uri, "outputs", self.video_id)
436                 os.makedirs(outputs, exist_ok=True)
437                 with open(os.path.join(outputs, "report.json"), "w", encoding="utf-8") as f:
438                     json.dump({"video_id": self.video_id, "status": "success"}, f)
439                 # deliberately DO NOT write intervals.csv
440                 os.makedirs(os.path.dirname(done_uri), exist_ok=True)
441                 with open(done_uri, "w", encoding="utf-8") as f:
442                     json.dump({"video_id": self.video_id, "returncode": 0}, f)
443                 return "exec"
444
445         _inject_fake(monkeypatch, _NoCsvClient())
446         out = m._execute_processing(
447             m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
448         )
449         assert out["status"] == "failed"
450         assert "could not read the result" in out["message"]
451         assert db.query_play_segments(out["video_id"], {}) == []
452
453
454         def test_cloud_video_id_divergence_is_tolerated_with_warning(cloud_env, caplog):
455             """If the worker re-hashes different bytes (result.video_id != local id), ingest
still proceeds
456             under the LOCAL id but logs a WARNING so the orphaned-report condition is
diagnosable."""
457             db, _bucket, monkeypatch = cloud_env
458             _inject_fake(monkeypatch, _FakeCloudClient(video_id="DIVERGED"))
459             with caplog.at_level("WARNING", logger="backend.main"):
460                 out = m._execute_processing(
461                     m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="wasb_hybrid")
462                 )
463                 assert out["status"] == "success"
464                 assert len(db.query_play_segments(out["video_id"], {})) == 2
465                 assert any(
466                     "DIVERGED" in r.getMessage() and "!= local" in r.getMessage()
467                     for r in caplog.records
468                 )
469
470
471         class _InProcSeg:
472             """Minimal in-process segmenter (writes one play row) ? proves the local path
actually ran."""
473
474             def __init__(self, db_, cfg_):
475                 self.db = db_
476
477             def process_video(self, vid, path, *a, **k):
478                 sid = self.db.add_play_segment(vid, 0.0, 5.0)
479                 return [
480                     {
481                         "id": sid,
482                         "start_time": 0.0,
483                         "end_time": 5.0,
484                         "duration": 5.0,
485                         "metadata": {},
486                     }

```

```

487         ]
488
489
490     # --- item 3: SPA compute-picker (per-request compute override + cloud_available
491     capability) ---
492
493     def test_per_request_local_forces_in_process(cloud_env):
494         """Server configured for cloud (compute_target=cloud_run), but a per-request
495         compute='local'
496         forces the in-process segmenter and NEVER dispatches to Cloud Run."""
497         db, _bucket, monkeypatch = cloud_env
498         fake = _FakeCloudClient()
499         _inject_fake(monkeypatch, fake)
500         monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _InProcSeg)
501
502         out = m._execute_processing(
503             m.global_config,
504             m.db_instance,
505             m.ProcessRequest(
506                 video_path="gs://b/clip.mp4", segmenter="motion", compute="local"
507             ),
508         )
509         assert out["status"] == "success"
510         assert len(out["play_segments"]) == 1
511         assert "cloud_run" not in out["message"]
512         assert (
513             fake.calls == []
514         ) # explicit "run on my machine" ? no dispatch despite a cloud server
515
516     def test_per_request_cloud_forces_dispatch_on_local_server(tmp_path, monkeypatch):
517         """Server default is LOCAL (compute_target=''), but a per-request compute='cloud'
518         forces a Cloud
519         Run dispatch when the control plane is wired (CLOUD_RUN_JOB + GOOGLE_CLOUD_PROJECT +
520         a bucket)."""
521         db = Database(str(tmp_path / "t.db")) # noqa: F841
522         monkeypatch.setattr(m, "db_instance", db)
523         bucket = tmp_path / "bucket"
524         bucket.mkdir()
525         cfg = AppConfig.model_validate(
526             {
527                 # noqa: F841
528                 "deployment": {"compute_target": ""}, # server default = in-process
529                 "storage": {"backend": "local", "output_root": str(bucket)},
530             }
531         )
532         monkeypatch.setattr(m, "global_config", cfg)
533         monkeypatch.setattr(m, "_get_executor", lambda: _InlineExecutor())
534         fake_local = tmp_path / "fetched.mp4"
535         fake_local.write_bytes(b"FAKEVIDEOBYTES")
536         monkeypatch.setattr(
537             m.storage, "acquire_local_input", lambda path, scfg: str(fake_local)
538         )
539         monkeypatch.setattr(
540             m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {})
541         )
542         monkeypatch.setenv("CLOUD_RUN_JOB", "rally-worker")
543         monkeypatch.setenv("GOOGLE_CLOUD_PROJECT", "proj")
544         fake = _FakeCloudClient()
545         _inject_fake(monkeypatch, fake)
546
547         out = m._execute_processing(
548             m.global_config,
549             m.db_instance,
550             m.ProcessRequest(

```

```

548         video_path="gs://b/clip.mp4", segmenter="wasb_hybrid", compute="cloud"
549     ),
550 )
551 assert out["status"] == "success"
552 assert len(out["play_segments"]) == 2
553 assert (
554     len(fake.calls) == 1
555 ) # per-request override dispatched despite the LOCAL server default
556
557
558 def test_per_request_cloud_unconfigured_fails_clearly(tmp_path, monkeypatch):
559     """compute='cloud' on a server NOT wired for it fails clearly (no silent local
560     fallback, no
561     dispatch) so the picker can surface 'cloud isn't available here'."""
562     db = Database(str(tmp_path / "t.db")) # noqa: F841
563     monkeypatch.setattr(m, "db_instance", db)
564     cfg = AppConfig.model_validate(
565         {"deployment": {"compute_target": ""}}
566     ) # no bucket, no env # noqa: F841
567     monkeypatch.setattr(m, "global_config", cfg)
568     monkeypatch.setattr(
569         m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
570     )
571     monkeypatch.delenv("CLOUD_RUN_JOB", raising=False)
572     monkeypatch.delenv("GOOGLE_CLOUD_PROJECT", raising=False)
573     video = tmp_path / "v.mp4"
574     video.write_bytes(b"hashable-bytes")
575
576     out = m._execute_processing(
577         m.global_config,
578         m.db_instance,
579         m.ProcessRequest(
580             video_path=str(video), segmenter="wasb_hybrid", compute="cloud"
581         ),
582     )
583     assert out["status"] == "failed"
584     assert "isn't configured" in out["message"]
585
586
587 def test_unknown_compute_value_falls_back_to_server_default(cloud_env, caplog):
588     """An unrecognised compute value is ignored (logged) and the server default applies
589     ? here a
590     cloud_run server still dispatches to the cloud."""
591     db, _bucket, monkeypatch = cloud_env
592     fake = _FakeCloudClient()
593     _inject_fake(monkeypatch, fake)
594     with caplog.at_level("WARNING", logger="backend.main"):
595         out = m._execute_processing(
596             m.global_config,
597             m.db_instance,
598             m.ProcessRequest(
599                 video_path="gs://b/clip.mp4", segmenter="wasb_hybrid", compute="banana"
600             ),
601         )
602     assert out["status"] == "success"
603     assert len(fake.calls) == 1 # fell back to the configured cloud_run default
604     assert any("unknown compute" in r.getMessage() for r in caplog.records)
605
606
607 def test_capabilities_reports_cloud_available(cloud_env):
608     """A cloud_run-configured server advertises deployment.cloud_available=True (the SPA
609     shows the
610     cloud toggle only then)."""
611     _db, _bucket, _monkeypatch = cloud_env
612     app = m.create_app(m.global_config, m.db_instance)

```

```

610     client = TestClient(app)
611     resp = client.get("/api/capabilities")
612     caps = resp.json()
613     assert caps["deployment"]["cloud_available"] is True
614     assert caps["deployment"]["compute_target"] == "cloud_run"
615
616
617     def test_capabilities_cloud_unavailable_by_default(monkeypatch, tmp_path):
618         """A default (local) server with no Cloud Run env advertises
cloud_available=False."""
619         cfg = AppConfig()
620         db = Database(str(tmp_path / "t.db"))
621         app = m.create_app(cfg, db)
622         monkeypatch.delenv("CLOUD_RUN_JOB", raising=False)
623         monkeypatch.delenv("GOOGLE_CLOUD_PROJECT", raising=False)
624         client = TestClient(app)
625         resp = client.get("/api/capabilities")
626         caps = resp.json()
627         assert caps["deployment"]["cloud_available"] is False
628
629
630     def test_full_api_process_to_query_via_cloud(cloud_env):
631         """End-to-end through the public API: POST /api/process (gs:// input) ? the inline
worker dispatches
632         to (fake) Cloud Run ? job done ? POST /api/query returns the ingested segments."""
633         db, _bucket, monkeypatch = cloud_env
634         fake = _FakeCloudClient(video_id="vidCloud2")
635         _inject_fake(monkeypatch, fake)
636         app = m.create_app(m.global_config, db)
637         client = TestClient(app)
638
639         r = client.post(
640             "/api/process",
641             json={"video_path": "gs://b/clip.mp4", "segmenter": "wasb_hybrid"},
642         )
643         assert r.status_code == 202
644         job_id = r.json()["job_id"]
645
646         job = client.get(f"/api/jobs/{job_id}").json()
647         assert job["status"] == "done"
648         video_id = job["video_id"]
649         assert video_id
650
651         q = client.post("/api/query", json={"video_id": video_id})
652         assert q.status_code == 200
653         segs = q.json()["play_segments"]
654         assert len(segs) == 2
655         assert all(_meta(s)["source"] == "cloud_run" for s in segs)
656
657
658     def test_local_override_survives_api_boundary(cloud_env):
659         """REGRESSION: the picker choice must survive the async submit?worker?reconstruct
boundary. On a
660         cloud_run server, POST /api/process with compute='local' must persist (jobs.compute)
+ reconstruct
661         (_job_to_request) the choice and run IN-PROCESS ? not dispatch to Cloud Run. (An
earlier cut
662         dropped compute in create_job/_job_to_request, silently nullifying the picker
through /api/process.)"""
663         db, _bucket, monkeypatch = cloud_env
664         fake = _FakeCloudClient()
665         _inject_fake(monkeypatch, fake)
666         monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _InProcSeg)
667         app = m.create_app(m.global_config, db)
668         client = TestClient(app)

```

```

669
670     r = client.post(
671         "/api/process",
672         json={
673             "video_path": "gs://b/clip.mp4",
674             "segmenter": "motion",
675             "compute": "local",
676         },
677     )
678     assert r.status_code == 202
679     job = client.get(f"/api/jobs/{r.json()['job_id']}").json()
680     assert job["status"] == "done"
681     assert (
682         fake.calls == []
683     ) # 'local' honoured across the boundary ? never dispatched to Cloud Run
684     segs = db.query_play_segments(job["video_id"], {})
685     assert len(segs) == 1
686     assert all(
687         _meta(s).get("source") != "cloud_run" for s in segs
688     ) # in-process rows, not cloud
689
690
691     def test_cloud_override_survives_api_boundary_on_local_server(tmp_path, monkeypatch):
692         """REGRESSION (mirror): on a LOCAL-default server wired for cloud (env + bucket),
693 POST /api/process
694         with compute='cloud' must survive the boundary and dispatch to Cloud Run."""
695         db = Database(str(tmp_path / "t.db")) # noqa: F841
696         monkeypatch.setattr(m, "db_instance", db)
697         bucket = tmp_path / "bucket"
698         bucket.mkdir()
699         cfg = AppConfig.model_validate(
700             { # noqa: F841
701                 "deployment": {"compute_target": ""}, # server default = in-process
702                 "storage": {"backend": "local", "output_root": str(bucket)},
703             }
704         )
705         monkeypatch.setattr(m, "global_config", cfg)
706         monkeypatch.setattr(m, "_get_executor", lambda: _InlineExecutor())
707         fake_local = tmp_path / "fetched.mp4"
708         fake_local.write_bytes(b"FAKEVIDEOBYTES")
709         monkeypatch.setattr(
710             m.storage, "acquire_local_input", lambda path, scfg: str(fake_local)
711         )
712         monkeypatch.setattr(
713             m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {})
714         )
715         monkeypatch.setenv("CLOUD_RUN_JOB", "rally-worker")
716         monkeypatch.setenv("GOOGLE_CLOUD_PROJECT", "proj")
717         fake = _FakeCloudClient(video_id="vidBoundary")
718         _inject_fake(monkeypatch, fake)
719         app = m.create_app(m.global_config, db)
720         client = TestClient(app)
721
722         r = client.post(
723             "/api/process",
724             json={
725                 "video_path": "gs://b/clip.mp4",
726                 "segmenter": "wasb_hybrid",
727                 "compute": "cloud",
728             },
729         )
730         assert r.status_code == 202
731         job = client.get(f"/api/jobs/{r.json()['job_id']}").json()
732         assert job["status"] == "done"
733         assert (

```

```

733         len(fake.calls) == 1
734     ) # 'cloud' honoured across the boundary ? dispatched despite local default
735     segs = db.query_play_segments(job["video_id"], {})
736     assert len(segs) == 2
737     assert all(_meta(s)["source"] == "cloud_run" for s in segs)
738
739
740     # --- hosted control-plane guardrail: compute='local' for a weights-backed detector
741     # -----
742     # A hosted cloud-serving control-plane bakes NO detector weights (only the worker Cloud
743     # Run Job mounts
744     # them), so a forced compute='local' for wasb/tracknet/fusion can't run in-process
745     # there. It must fail
746     # FAST with an actionable "run in the cloud" message ? not fall through to the in-
747     # process segmenter
748     # that dies deep with a cryptic "WASB weights not found" (observed live: job f29d41d0?,
749     # 2026-07-07).
750
751     def test_local_on_cloud_server_without_weights_fails_fast(cloud_env, monkeypatch):
752         """compute='local' + a weights-backed detector on a cloud_run server with NO local
753         weights ?
754         a clear pre-dispatch failure (never dispatched, never the cryptic deep 'weights not
755         found')."""
756         _db, _bucket, _mp = cloud_env
757         monkeypatch.delenv("WASB_WEIGHTS", raising=False)
758         m.global_config.indexing.wasb.weights_path =
759         "/nonexistent/wasb_badminton_best.pth.tar"
760         fake = _FakeCloudClient()
761         _inject_fake(monkeypatch, fake)
762
763         out = m._execute_processing(
764             m.ProcessRequest(
765                 video_path="gs://b/clip.mp4", segmenter="wasb_hybrid", compute="local"
766             )
767         )
768         assert out["status"] == "failed"
769         assert "can't run detection locally" in out["message"]
770         assert "Run in cloud" in out["message"]
771         assert "weights not found" not in out["message"] # NOT the cryptic deep failure
772         assert fake.calls == [] # never dispatched
773         # Honest provenance: requested local, ran NOWHERE (pre-dispatch reject).
774         assert out["compute"]["requested"] == "local"
775         assert out["compute"]["dispatched"] is False
776         assert out["compute"]["path"] == "not_run"
777
778     def test_local_on_cloud_server_with_weights_present_runs_in_process(
779         cloud_env, tmp_path, monkeypatch
780     ):
781         """Escape hatch: the guardrail must NOT over-block ? on a cloud-wired box that
782         ACTUALLY has the
783         detector weights on disk, compute='local' still runs the in-process segmenter (no
784         dispatch)."""
785         _db, _bucket, _mp = cloud_env
786         weights = tmp_path / "wasb_badminton_best.pth.tar"
787         weights.write_bytes(b"WEIGHTS")
788         monkeypatch.delenv("WASB_WEIGHTS", raising=False)
789         m.global_config.indexing.wasb.weights_path = str(weights)
790         monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _InProcSeg)
791         fake = _FakeCloudClient()
792         _inject_fake(monkeypatch, fake)
793
794         out = m._execute_processing(
795             m.ProcessRequest(

```



```

788         video_path="gs://b/clip.mp4", segmenter="wasb_hybrid", compute="local"
789     )
790 )
791 assert out["status"] == "success"
792 assert fake.calls == [] # in-process, not dispatched
793 assert out["compute"]["path"] == "in_process"
794
795
796 def test_local_non_weights_segmenter_on_cloud_server_still_runs_in_process(
797     cloud_env, monkeypatch
798 ):
799     """The guardrail is scoped to weights-backed detectors: a weight-free segmenter
800 (motion) with
801     compute='local' on a cloud server is unaffected ? it runs in-process, never fails
802 fast."""
803     _db, _bucket, _mp = cloud_env
804     m.global_config.indexing.wasb.weights_path =
805 "/nonexistent/should_not_matter.pth.tar"
806     monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _InProcSeg)
807     fake = _FakeCloudClient()
808     _inject_fake(monkeypatch, fake)
809
810     out = m._execute_processing(
811         m.ProcessRequest(
812             video_path="gs://b/clip.mp4", segmenter="motion", compute="local"
813         )
814     )
815     assert out["status"] == "success"
816     assert fake.calls == []
817     assert out["compute"]["path"] == "in_process"
818
819
820 def test_local_weights_backed_on_pure_local_server_is_not_guardrailed(
821     tmp_path, monkeypatch
822 ):
823     """The guardrail only fires on a cloud-capable server. On a pure-local server (no
824 cloud), a
825     weights-backed compute='local' request falls through to the in-process segmenter as
826 before ? the
827     'run in the cloud' message would be nonsense there (no cloud to fall back to)."""
828     db = Database(str(tmp_path / "t.db")) # noqa: F841
829     monkeypatch.setattr(m, "db_instance", db)
830     cfg = AppConfig.model_validate({"deployment": {"compute_target": ""}}) # noqa: F841
831     monkeypatch.setattr(m, "global_config", cfg)
832     monkeypatch.setattr(
833         m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {})
834     )
835     monkeypatch.delenv("CLOUD_RUN_JOB", raising=False)
836     monkeypatch.delenv("GOOGLE_CLOUD_PROJECT", raising=False)
837     monkeypatch.delenv("WASB_WEIGHTS", raising=False)
838     cfg.indexing.wasb.weights_path = "/nonexistent/wasb.pth.tar"
839     monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _InProcSeg)
840     video = tmp_path / "v.mp4"
841     video.write_bytes(b"hashable-bytes")
842
843     out = m._execute_processing(
844         m.ProcessRequest(
845             video_path=str(video), segmenter="wasb_hybrid", compute="local"
846         )
847     )
848     # Not the hosted fast-fail ? it ran in-process (the local box owns its own weights
849 story).
850     assert out["status"] == "success"
851     assert out["compute"]["path"] == "in_process"
852

```

```

847
848     def test_local_tracknet_on_cloud_server_without_weights_fails_fast(cloud_env,
monkeypatch):
849         """The guardrail also covers tracknet_hybrid (not just wasb): with no TrackNet
weights configured
850         (the default), compute='local' on a cloud_run server fails fast, never
dispatching."""
851         _db, _bucket, _mp = cloud_env
852         monkeypatch.delenv("WASB_WEIGHTS", raising=False)
853         m.global_config.indexing.tracknet.weights_path = "" # default: no local TrackNet
weights
854         fake = _FakeCloudClient()
855         _inject_fake(monkeypatch, fake)
856
857         out = m._execute_processing(
858             m.ProcessRequest(
859                 video_path="gs://b/clip.mp4", segmenter="tracknet_hybrid", compute="local"
860             )
861         )
862         assert out["status"] == "failed"
863         assert "can't run detection locally" in out["message"]
864         assert fake.calls == []
865         assert out["compute"]["path"] == "not_run"
866
867
868     def test_local_fusion_tracknet_substrate_checks_tracknet_not_wasb(
869         cloud_env, tmp_path, monkeypatch
870     ):
871         """Pins the substrate switch: fusion_hybrid on the TrackNet substrate must gate on
the TrackNet
872         weights, not WASB. WASB weights PRESENT but TrackNet weights ABSENT must still fail
fast ? if the
873         check wrongly read indexing.wasb.weights_path it would take the escape hatch and run
in-process."""
874         _db, _bucket, _mp = cloud_env
875         wasb = tmp_path / "wasb_badminton_best.pth.tar"
876         wasb.write_bytes(b"WEIGHTS")
877         monkeypatch.delenv("WASB_WEIGHTS", raising=False)
878         m.global_config.indexing.wasb.weights_path = str(wasb) # WASB present ?
879         m.global_config.indexing.fusion.substrate = "tracknet"
880         m.global_config.indexing.tracknet.weights_path = "/nonexistent/tracknet.pth" # ?
TrackNet absent
881         fake = _FakeCloudClient()
882         _inject_fake(monkeypatch, fake)
883
884         out = m._execute_processing(
885             m.ProcessRequest(
886                 video_path="gs://b/clip.mp4", segmenter="fusion_hybrid", compute="local"
887             )
888         )
889         assert out["status"] == "failed" # gated on TrackNet weights, which are absent
890         assert "can't run detection locally" in out["message"]
891         assert fake.calls == []
892
893
894     def test_local_fusion_default_substrate_with_wasb_weights_runs_in_process(
895         cloud_env, tmp_path, monkeypatch
896     ):
897         """Escape hatch for fusion on its DEFAULT (wasb) substrate: with WASB weights
present, compute=
898         'local' runs in-process (proves fusion_hybrid reads wasb weights on the default
substrate)."""
899         _db, _bucket, _mp = cloud_env
900         wasb = tmp_path / "wasb_badminton_best.pth.tar"
901         wasb.write_bytes(b"WEIGHTS")

```

```

902 monkeypatch.delenv("WASB_WEIGHTS", raising=False)
903 m.global_config.indexing.wasb.weights_path = str(wasb)
904 # fusion.substrate defaults to "wasb"
905 monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _InProcSeg)
906 fake = _FakeCloudClient()
907 _inject_fake(monkeypatch, fake)
908
909 out = m._execute_processing(
910     m.ProcessRequest(
911         video_path="gs://b/clip.mp4", segmenter="fusion_hybrid", compute="local"
912     )
913 )
914 assert out["status"] == "success"
915 assert fake.calls == []
916 assert out["compute"]["path"] == "in_process"
917
918
919 # --- compute-path PROVENANCE (the validation method: prove which backend ran)
-----
920
921
922 def test_cloud_dispatch_stamps_compute_provenance(cloud_env):
923     """A cloud run stamps result['compute'] with path=cloud_run + the REAL Cloud Run
execution name
924     (the GCP-verifiable proof) + the requested choice + job/region/outputs ? so the
intended path is
925     provable from the job result alone (no guessing from segment metadata)."""
926     _db, _bucket, monkeypatch = cloud_env
927     _inject_fake(monkeypatch, _FakeCloudClient())
928
929     out = m._execute_processing(
930         m.ProcessRequest(
931             video_path="gs://b/clip.mp4", segmenter="wasb_hybrid", compute="cloud"
932         )
933     )
934     prov = out["compute"]
935     assert prov["path"] == "cloud_run"
936     assert prov["requested"] == "cloud"
937     assert (
938         prov["execution"] == "exec-1"
939     ) # the fake client's execution name (real one in prod)
940     assert prov["outputs_uri"].endswith("/outputs/vidCloud/")
941     assert prov["job"] and prov["region"] # the dispatch coordinates are recorded
942
943
944 def test_local_override_stamps_in_process_provenance(cloud_env):
945     """On a cloud server, compute='local' runs in-process AND the result proves it:
path=in_process,
946     requested=local ? so a 'ran local despite a cloud server' is auditable (no silent
ambiguity)."""
947     _db, _bucket, monkeypatch = cloud_env
948     _inject_fake(monkeypatch, _FakeCloudClient())
949     monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _InProcSeg)
950
951     out = m._execute_processing(
952         m.ProcessRequest(
953             video_path="gs://b/clip.mp4", segmenter="motion", compute="local"
954         )
955     )
956     assert out["status"] == "success"
957     assert out["compute"] == {"path": "in_process", "requested": "local"}
958
959
960 def test_default_in_process_stamps_provenance(tmp_path, monkeypatch):
961     """A default-OFF (no cloud) run stamps path=in_process, requested=None ? the byte-

```

```

identical local
962     path is now self-describing."""
963     db = Database(str(tmp_path / "t.db")) # noqa: F841
964     monkeypatch.setattr(m, "db_instance", db)
965     monkeypatch.setattr(m, "db_instance", db)
966     monkeypatch.setattr(m, "global_config", AppConfig())
967     m.app.state.config = AppConfig()
968     monkeypatch.setattr(
969         m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
970 )
971 monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _InProcSeg)
972 video = tmp_path / "v.mp4"
973 video.write_bytes(b"hashable-bytes")
974
975 out = m._execute_processing(
976     m.ProcessRequest(video_path=str(video), segmenter="motion")
977 )
978 assert out["compute"] == {"path": "in_process", "requested": None}
979
980
981 def test_cloud_requested_unconfigured_stamps_not_dispatched(tmp_path, monkeypatch):
982     """compute='cloud' on a server that can't dispatch fails clearly AND records it:
target cloud_run,
983     requested=cloud, dispatched=False ? proving the cloud job never ran (no silent local
fallback)."""
984     db = Database(str(tmp_path / "t.db")) # noqa: F841
985     monkeypatch.setattr(m, "db_instance", db)
986     monkeypatch.setattr(m, "db_instance", db)
987     monkeypatch.setattr(
988         m,
989         "global_config",
990         AppConfig.model_validate({"deployment": {"compute_target": ""}}),
991     )
992     m.app.state.config = AppConfig.model_validate(
993         {"deployment": {"compute_target": ""}}
994     )
995     monkeypatch.setattr(
996         m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
997 )
998 monkeypatch.delenv("CLOUD_RUN_JOB", raising=False)
999 monkeypatch.delenv("GOOGLE_CLOUD_PROJECT", raising=False)
1000 video = tmp_path / "v.mp4"
1001 video.write_bytes(b"hashable-bytes")
1002
1003 out = m._execute_processing(
1004     m.ProcessRequest(
1005         video_path=str(video), segmenter="wasb_hybrid", compute="cloud"
1006     )
1007 )
1008 assert out["status"] == "failed"
1009 # Honest: it ran NOWHERE (pre-dispatch reject) ? path=not_run, but the intent is
recorded.
1010 assert out["compute"] == {
1011     "path": "not_run",
1012     "requested": "cloud",
1013     "target": "cloud_run",
1014     "dispatched": False,
1015 }
1016
1017
1018 def test_validation_failure_stamps_not_run_provenance(tmp_path, monkeypatch):
1019     """A validation failure never reaches a segmenter ? path=not_run (nothing
executed)."""
1020     db = Database(str(tmp_path / "t.db")) # noqa: F841
1021     monkeypatch.setattr(m, "db_instance", db)

```

```

1022 monkeypatch.setattr(m, "global_config", AppConfig())
1023 m.app.state.config = AppConfig()
1024 fail = [
1025     ValidationResult(validator_name="x", passed=False, message="bad", details={})
1026 ]
1027 monkeypatch.setattr(m.registry, "run_all_with_context", lambda *a, **k: (fail, {}))
1028 video = tmp_path / "v.mp4"
1029 video.write_bytes(b"hashable-bytes")
1030
1031 out = m._execute_processing(m.ProcessRequest(video_path=str(video)))
1032 assert out["status"] == "failed"
1033 assert out["compute"]["path"] == "not_run"
1034
1035
1036 def test_full_api_process_job_result_carries_cloud_provenance(cloud_env):
1037     """End-to-end: POST /api/process (compute=cloud) ? the job result the SPA polls
carries the
1038     cloud_run provenance + execution, so the UI/operator can PROVE the cloud path was
taken."""
1039     _db, _bucket, monkeypatch = cloud_env
1040     _inject_fake(monkeypatch, _FakeCloudClient(video_id="vidProv"))
1041     app = m.create_app(m.global_config, _db)
1042     client = TestClient(app)
1043
1044     r = client.post(
1045         "/api/process",
1046         json={
1047             "video_path": "gs://b/clip.mp4",
1048             "segmenter": "wasb_hybrid",
1049             "compute": "cloud",
1050         },
1051     )
1052     assert r.status_code == 202
1053     job = client.get(f"/api/jobs/{r.json()['job_id']}").json()
1054     assert job["status"] == "done"
1055     assert job["result"]["compute"]["path"] == "cloud_run"
1056     assert job["result"]["compute"]["execution"] == "exec-1"
1057
1058
1059 # --- #468: local compute lane ? CPU/GPU-heavy in-process work serializes, remote polls
don't -----
1060
1061
1062 def test_in_process_detection_runs_under_the_local_lane(cloud_env):
1063     """The in-process segmenter holds orchestration._LOCAL_COMPUTE_LANE while it runs,
so raising the
1064     drain worker count (to overlap I/O-bound remote dispatch) never runs two local
detections at once.
1065     A remote dispatch returns before this path, so remote polls never hold the lane
(they overlap)."""
1066     import backend.orchestration as orch
1067
1068     _db, _bucket, monkeypatch = cloud_env
1069     held = {}
1070
1071     class _LaneCheckSeg:
1072         def __init__(self, db_, cfg_):
1073             self.db = db_
1074
1075         def process_video(self, vid, path, *a, **k):
1076             # the wrapping `with _LOCAL_COMPUTE_LANE` holds the size-1 lane ? a non-
blocking re-acquire
1077             # must FAIL while detection runs (proving concurrent local detections are
excluded).
1078             got = orch._LOCAL_COMPUTE_LANE.acquire(blocking=False)

```

```

1079         held["busy"] = not got
1080         if got:
1081             orch._LOCAL_COMPUTE_LANE.release()
1082         sid = self.db.add_play_segment(vid, 0.0, 5.0)
1083         return [
1084             {"id": sid, "start_time": 0.0, "end_time": 5.0, "duration": 5.0,
"metadata": {}}
1085         ]
1086
1087         monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _LaneCheckSeg)
1088         out = m._execute_processing(
1089             m.global_config,
1090             m.db_instance,
1091             m.ProcessRequest(video_path="gs://b/clip.mp4", segmenter="motion",
compute="local"),
1092         )
1093         assert out["status"] == "success"
1094         assert held.get("busy") is True # detection ran while the local lane was held
1095
1096
1097     def test_wait_progress_message_is_honest_and_minute_granular():
1098         """#482: the CP knows ELAPSED time, not frames ? the heartbeat says so plainly,
shows
1099         movement from the first tick ("starting up" during queue/cold-start), then counts
minutes."""
1100         from backend.orchestration import _wait_progress_message
1101
1102         early = _wait_progress_message("wasb_hybrid", 12.0)
1103         assert "wasb_hybrid" in early and "starting up" in early
1104         later = _wait_progress_message("wasb_hybrid", 125.0)
1105         assert later.endswith("2 min elapsed")
1106

```

38. assistant (2026-07-08T05:12:04.778Z)

I have a thorough understanding. Let me set up the isolated worktree (per branching-safety rules) and check the test environment, while reading the remaining test patterns.

39. user (2026-07-08T05:12:09.087Z)

```

=== venv? ===
=== python ===
Python 3.13.13
=== ruff/mypy ===
ruff 0.15.16
mypy 2.1.0 (compiled: yes)
=== existing worktrees ===
C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer
699d85b [master]
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/cfg-wt      20b2d5a [claude/wasb-
config-validation]
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/compile-wt  b09a1d4 [claude/compile-
reliability]
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/pr458-pre-wt 355dad6 (detached HEAD)
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/6e9d0958-a2ec-4a8e-bb2d-6cce9833c527/scratchpad/pr458-wt    5de1ef2 (detached HEAD)

```

40. user (2026-07-08T05:12:09.105Z)

```

1         """Worker dispatcher: pull (storage) ? run pipeline ? push (storage).
2

```

```

3     `cmd_infer` is exercised end-to-end over the LOCAL storage backend with the validator +
4     segmenter mocked, so it needs no real video/GPU/network. The real S3 round-trip is
validated
5     on the droplet against MinIO.
6     """
7
8     import json
9
10    import pytest
11
12    from backend.config import load_config
13    from backend.results import Failure
14    from backend.worker import __main__ as wmain
15
16    # ----- pure helpers
17
18
19    @pytest.mark.parametrize(
20        "raw,expected",
21        [
22            ("true", True),
23            ("False", False),
24            ("3", 3),
25            ("2.5", 2.5),
26            ("wasb_hybrid", "wasb_hybrid"),
27        ],
28    )
29    def test_coerce(raw, expected):
30        assert wmain._coerce(raw) == expected and type(wmain._coerce(raw)) is type(expected)
31
32
33    def test_apply_overrides_on_typed_config():
34        cfg = _fresh_config()
35        wmain._apply_overrides(
36            cfg,
37            [
38                "indexing.skip_ai_handoff=true",
39                "indexing.min_segment_duration=1.5",
40                "sport=tennis",
41            ],
42        )
43        assert cfg.indexing.skip_ai_handoff is True
44        assert cfg.indexing.min_segment_duration == 1.5
45        assert cfg.sport == "tennis"
46
47
48    def test_apply_overrides_rejects_bad_format():
49        cfg = _fresh_config()
50        with pytest.raises(ValueError):
51            wmain._apply_overrides(cfg, ["no_equals_sign"])
52
53
54    def test_write_intervals_csv(tmp_path):
55        segs = [
56            {
57                "start_time": 2.0,
58                "end_time": 5.0,
59                "shot_count": 4,
60                "shot_count_confidence": "LocalNoAI",
61            },
62            {"start_time": 9.0, "end_time": 12.5},
63        ]
64        out = tmp_path / "intervals.csv"
65        wmain._write_intervals_csv(segs, str(out))
66        lines = out.read_text().splitlines()

```

```

67         assert lines[0] == "start,end,shot_count,ending_reason"
68         assert lines[1].startswith("2.000,5.000,4,")
69         assert lines[2].startswith("9.000,12.500,,")
70
71
72     def test_write_report_json(tmp_path):
73         out = tmp_path / "report.json"
74         wmain._write_report_json(
75             "vid1", [{"start_time": 1.0, "end_time": 2.0}], "success", str(out)
76         )
77         data = json.loads(out.read_text())
78         assert data["video_id"] == "vid1" and data["segment_count"] == 1
79         assert data["segments"][0] == {"start": 1.0, "end": 2.0}
80
81
82     def test_parser_infer_accepts_set_and_uris():
83         args = wmain.build_parser().parse_args(
84             [
85                 "infer",
86                 "--video-uri",
87                 "s3://b/v.mp4",
88                 "--out-uri",
89                 "s3://b",
90                 "--set",
91                 "indexing.skip_ai_handoff=true",
92                 "--set",
93                 "sport=tennis",
94             ]
95         )
96         assert args.cmd == "infer" and args.video_uri == "s3://b/v.mp4"
97         assert args.set == ["indexing.skip_ai_handoff=true", "sport=tennis"]
98
99
100     def test_parser_infer_skip_validation_flag():
101         base = ["infer", "--video-uri", "s3://b/v.mp4", "--out-uri", "s3://b"]
102         # Default-OFF: the direct `worker infer` CLI still validates.
103         assert wmain.build_parser().parse_args(base).skip_validation is False
104         # The dispatcher passes it to disable the redundant worker re-validation.
105         assert wmain.build_parser().parse_args(base + ["--skip-validation"]).skip_validation
106         is True
107
108     # ----- cmd_infer (local,
109     mocked)
110
111     class _OKResult:
112         passed = True
113         validator_name = "fake"
114         message = "ok"
115
116         def model_dump(self):
117             return {"validator_name": "fake", "passed": True, "message": "ok"}
118
119
120     class _FailResult(_OKResult):

```

41. assistant (2026-07-08T05:12:22.907Z)

No repo venv ? deps are in global Python 3.13, so a worktree anywhere can run the suite. There are sibling-session worktrees on other branches (the WIP the task warned about) ? I'll stay isolated. Let me create my worktree from `origin/master`.

42. user (2026-07-08T05:12:26.505Z)


```
Exit code 1
Preparing worktree (new branch 'feat/521-compile-on-l4')
branch 'feat/521-compile-on-l4' set up to track 'origin/master'.
error: unable to create file docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-
serving_gpu-resident-nvdec_2026-07-07.md: Filename too long
error: unable to create file docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-
serving_validation-latency_2026-07-07.md: Filename too long
fatal: Could not reset index file to revision 'HEAD'.
---VERIFY---
/usr/bin/bash: line 1: cd: C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-
khelsutra-guru-badminton-highlight-
indexer/75d3d121-f240-4293-afa5-b436a379ad31/scratchpad/wt-521: No such file or directory
```

43. assistant (2026-07-08T05:12:33.414Z)

Windows MAX_PATH hit on the long scratchpad path. Let me clean up and use a short worktree path.

44. user (2026-07-08T05:12:41.620Z)

```
error: failed to delete '.git/worktrees/a1-master-worktree': Permission denied
error: failed to delete '.git/worktrees/bhi-wt-worker-reval': Permission denied
error: failed to delete '.git/worktrees/fix-worktree': Permission denied
error: failed to delete '.git/worktrees/master-compare': Permission denied
error: failed to delete '.git/worktrees/pr465': Permission denied
error: failed to delete '.git/worktrees/pr465-current-merge': Permission denied
error: failed to delete '.git/worktrees/pr465-merge': Permission denied
error: failed to delete '.git/worktrees/pr467': Permission denied
error: failed to delete '.git/worktrees/pr469': Permission denied
error: failed to delete '.git/worktrees/wt-465fix': Permission denied
error: failed to delete '.git/worktrees/wt-a1-runlog': Permission denied
error: failed to delete '.git/worktrees/wt-claudemd-r25': Permission denied
error: failed to delete '.git/worktrees/wt-cloudrun-cancel': Permission denied
error: failed to delete '.git/worktrees/wt-cloudrun-dispatch': Permission denied
error: failed to delete '.git/worktrees/wt-cp-rebuild-spec': Permission denied
error: failed to delete '.git/worktrees/wt-dockerfile-auth': Permission denied
error: failed to delete '.git/worktrees/wt-docs': Permission denied
error: failed to delete '.git/worktrees/wt-docs-reconcile': Permission denied
error: failed to delete '.git/worktrees/wt-drain-concurrency': Permission denied
error: failed to delete '.git/worktrees/wt-ffmpeg-timeouts': Permission denied
error: failed to delete '.git/worktrees/wt-gdrive-api-safety': Permission denied
error: failed to delete '.git/worktrees/wt-idempotent-process': Permission denied
error: failed to delete '.git/worktrees/wt-input-jail': Permission denied
error: failed to delete '.git/worktrees/wt-output-dir-config': Permission denied
error: failed to delete '.git/worktrees/wt-p1-genyaml': Permission denied
error: failed to delete '.git/worktrees/wt-p2-upload-stream': Permission denied
error: failed to delete '.git/worktrees/wt-p4-md5cache': Permission denied
error: failed to delete '.git/worktrees/wt-p5a-pollknob': Permission denied
error: failed to delete '.git/worktrees/wt-p6-heartbeat': Permission denied
error: failed to delete '.git/worktrees/wt-pr481-review': Permission denied
error: failed to delete '.git/worktrees/wt-retention-sweep': Permission denied
error: failed to delete '.git/worktrees/wt-review-492-mergecheck': Permission denied
error: failed to delete '.git/worktrees/wt-review-493-mergecheck': Permission denied
error: failed to delete '.git/worktrees/wt-scratch-reclaim': Permission denied
error: failed to delete '.git/worktrees/wt-signed-download': Permission denied
error: failed to delete '.git/worktrees/wt-source-cache': Permission denied
error: failed to delete '.git/worktrees/wt-upload-dir': Permission denied
error: failed to delete '.git/worktrees/wt-upload-staging': Permission denied
Deleted branch feat/521-compile-on-l4 (was 699d85b).
--- retry with short path ---
Preparing worktree (new branch 'feat/521-compile-on-l4')
branch 'feat/521-compile-on-l4' set up to track 'origin/master'.
HEAD is now at 699d85b feat(validation): cache validation verdict by video_id to skip re-
validation on reupload (#518)
---VERIFY---
feat/521-compile-on-l4
```

```
699d85b feat(validation): cache validation verdict by video_id to skip re-validation on reupload
(#518)
---ONLY MY COMMITS (should be empty)---
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

45. assistant (2026-07-08T05:12:56.753Z)

Worktree is clean at `C:/kg-wt521` (off origin/master 699d85b, no sibling commits). Let me establish a green baseline before changing anything, and set up progress tracking.

46. user (2026-07-08T05:13:51.566Z)

```
..... [ 41%]
..... [ 83%]
..... [100%]
===== warnings summary =====
tests/test_api_cloud_dispatch.py::test_full_api_process_to_query_via_cloud
tests/test_api_cloud_dispatch.py::test_full_api_process_to_query_via_cloud
tests/test_api_cloud_dispatch.py::test_local_override_survives_api_boundary
tests/test_api_cloud_dispatch.py::test_local_override_survives_api_boundary
tests/test_api_cloud_dispatch.py::test_cloud_override_survives_api_boundary_on_local_server
tests/test_api_cloud_dispatch.py::test_cloud_override_survives_api_boundary_on_local_server
tests/test_api_cloud_dispatch.py::test_full_api_process_job_result_carries_cloud_provenance
tests/test_api_cloud_dispatch.py::test_full_api_process_job_result_carries_cloud_provenance
  C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\site-
packages\google\auth\_default.py:113: UserWarning: Your application has authenticated using end
user credentials from Google Cloud SDK without a quota project. You might receive a "quota
exceeded" or "API not enabled" error. See the following page for troubleshooting:
https://cloud.google.com/docs/authentication/adc-troubleshooting/user-creds.
    warnings.warn(_CLOUD_SDK_CREDENTIALS_WARNING)

-- Docs: https://docs.pytest.org/en/stable/how-to/capture-warnings.html
172 passed, 8 warnings in 50.78s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

47. assistant (2026-07-08T05:13:57.689Z)

Baseline is green (172 passed). Note: shell cwd resets to main repo each call, so I'll always `cd` to the worktree. Let me read the compute-backend and async-jobs test patterns, and the rest of test_worker.py.

48. user (2026-07-08T05:13:58.323Z)

```
1      """ComputeBackend seam ? Cloud Run dispatch shim (COMPUTE_DECOUPLED_SERVING M6).
2
3      All GPU-free + cloud-free: the Cloud Run trigger is a FAKE ``CloudRunJobClient`` that
simulates the
4      worker (writes the done-marker + report to a local "bucket" = a tmp dir), so the
dispatch?bucket-poll
5      ?read flow is proven for $0. The cardinal invariant: **default-OFF** ? only
``compute_target=cloud_run``
6      selects a backend; everything else returns ``None`` (= run the worker's existing in-
process path).
7      """
8
9      import json
10     import os
11     import tempfile
12
13     import pytest
14
15     from backend.compute import (
16         ComputeJobSpec,
17         RemoteExecutionFailed,
```

```

18         get_compute_backend,
19     )
20     from backend.compute.cloud_run import CloudRunCompute, CloudRunJobClient
21     from backend.config import AppConfig
22     from backend.infrastructure.execution_policy import ExecutionPolicy, PolicyTimeout
23
24     # A fast, no-wait policy (the fake client writes the marker synchronously, so the first
poll hits).
25     _FAST = ExecutionPolicy(poll_every_n_sec=0.0, op_timeout_sec=5.0, max_wait_sec=5.0)
26
27
28     class _FakeRunClient(CloudRunJobClient):
29         """Simulates a Cloud Run Job: on dispatch, write the report + done-marker into the
local bucket
30         exactly where the real worker would, so CloudRunCompute's bucket-poll resolves
immediately."""
31
32         def __init__(self, *, video_id="vidCloud", segment_count=2, returncode=0):
33             self.calls: list[dict] = []
34             self.video_id, self.segment_count, self.returncode = (
35                 video_id,
36                 segment_count,
37                 returncode,
38             )
39
40         def run_job(self, job_name, argv, *, region, project=None):
41             self.calls.append(
42                 {"job": job_name, "argv": list(argv), "region": region, "project": project}
43             )
44             out_uri = argv[argv.index("--out-uri") + 1].rstrip("/")
45             done_uri = argv[argv.index("--done-uri") + 1]
46             outputs = os.path.join(out_uri, "outputs", self.video_id)
47             os.makedirs(outputs, exist_ok=True)
48             with open(os.path.join(outputs, "report.json"), "w", encoding="utf-8") as f:
49                 json.dump(
50                     {
51                         "video_id": self.video_id,
52                         "segment_count": self.segment_count,
53                         "status": "success",
54                     },
55                     f,
56                 )
57             os.makedirs(os.path.dirname(done_uri), exist_ok=True)
58             with open(done_uri, "w", encoding="utf-8") as f:
59                 json.dump(
60                     {
61                         "video_id": self.video_id,
62                         "returncode": self.returncode,
63                         "segment_count": self.segment_count,
64                         "status": "success",
65                     },
66                     f,
67                 )
68             return "exec-123"
69
70         def cancel_execution(self, execution):
71             raise AssertionError(
72                 "cancel must never fire on the happy path"
73             ) # pragma: no cover
74
75
76     class _NoopRunClient(CloudRunJobClient):
77         """Records the dispatch but does NO filesystem writes ? for tests that fake the
storage layer
78         directly (a gs:// out_uri has no local dir to makedirs)."""

```

```

79
80     def __init__(self):
81         self.calls: list[dict] = []
82         self.cancelled: list[str] = []
83
84     def run_job(self, job_name, argv, *, region, project=None):
85         self.calls.append(
86             {"job": job_name, "argv": list(argv), "region": region, "project": project}
87         )
88         return "exec"
89
90     def cancel_execution(self, execution):
91         self.cancelled.append(execution)
92
93
94     class _NeverDoneClient(CloudRunJobClient):
95         """Simulates a still-running execution: dispatch succeeds (returning a fully-
qualified
96         execution name) but NO done-marker is ever written, so the bucket-poll must time
out."""
97
98         EXEC_NAME = "projects/p/locations/asia-south1/jobs/rally-worker/executions/exec-
slow"
99
100        def __init__(self, *, cancel_error: Exception | None = None):
101            self.calls: list[dict] = []
102            self.cancelled: list[str] = []
103            self._cancel_error = cancel_error
104
105        def run_job(self, job_name, argv, *, region, project=None):
106            self.calls.append(
107                {"job": job_name, "argv": list(argv), "region": region, "project": project}
108            )
109            return self.EXEC_NAME
110
111        def cancel_execution(self, execution):
112            self.cancelled.append(execution)
113            if self._cancel_error is not None:
114                raise self._cancel_error
115
116
117        class _CrashedNoMarkerClient(CloudRunJobClient):
118            """Simulates the bug's failure class: the worker execution TERMINATES (crash / OOM /
an old
119            image argparse-aborting on an unknown flag) WITHOUT ever writing a done-marker.
Dispatch returns
120            a resolvable execution name, no marker is ever written, and the Executions API
reports a terminal
121            state ? so run_infer must fail FAST via the status probe instead of spinning the
poll budget."""
122
123            EXEC_NAME = "projects/p/locations/asia-south1/jobs/rally-worker/executions/exec-
crash"
124
125            def __init__(self, *, state: str = "Failed"):
126                self.calls: list[dict] = []
127                self.cancelled: list[str] = []
128                self.status_polls = 0
129                self._state = state
130
131            def run_job(self, job_name, argv, *, region, project=None):
132                self.calls.append(
133                    {"job": job_name, "argv": list(argv), "region": region, "project": project}
134                )
135                return self.EXEC_NAME

```

```

136
137     def cancel_execution(self, execution):
138         self.cancelled.append(execution)
139
140     def execution_terminal_state(self, execution):
141         self.status_polls += 1
142         # Anti-hang guard: the fast path MUST abort on the first terminal observation.
143         # regression kept polling despite a terminal status, trip here instead of
144         # spinning.
145         assert self.status_polls <= 5, "status probed repeatedly ? fast path did not
146         abort the poll"
147         return self._state
148
149     # ----- factory (default-
150     OFF)
151     def _cfg(compute_target="", storage_backend="local"):
152         return AppConfig.model_validate(
153             {
154                 "deployment": {"compute_target": compute_target},
155                 "storage": {"backend": storage_backend},
156             }
157         )
158
159     def test_factory_is_none_for_local_targets():
160         """DEFAULT-OFF: no/local/cpu_mock compute target ? None (run the in-process
161         path)."""
162         assert get_compute_backend(_cfg("")) is None
163         assert get_compute_backend(_cfg("local_gpu")) is None
164         assert get_compute_backend(_cfg("cpu_mock")) is None
165
166     def test_factory_returns_cloud_run_for_cloud_target():
167         backend = get_compute_backend(_cfg("cloud_run", "gcs"), run_client=_FakeRunClient())
168         assert isinstance(backend, CloudRunCompute)
169
170     # ----- real client: missing-project
171     guard
172
173     def test_real_client_rejects_missing_project():
174         """The real GoogleCloudRunJobClient must fail loudly when GOOGLE_CLOUD_PROJECT is
175         unset ?
176         a None/empty project would build a bogus "projects/None/..." job path (job_path
177         requires a
178         real str). The guard runs BEFORE the google-cloud-run import, so it is exercisable
179         in CI
180         (the SDK is intentionally not a test dependency)."""
181         from backend.compute.cloud_run import GoogleCloudRunJobClient
182
183         client = GoogleCloudRunJobClient()
184         for missing in (None, ""):
185             with pytest.raises(RuntimeError, match="GOOGLE_CLOUD_PROJECT"):
186                 client.run_job(
187                     "rally-worker", ["infer"], region="asia-south1", project=missing
188                 )
189
190     def test_real_client_cancel_rejects_unresolvable_execution_name():
191         """cancel_execution must fail loudly on a non-resolvable execution name ? run_job

```

```

falls back
192         to the literal "execution" when the operation metadata carries no name, and
cancelling that
193         would be a bogus API call. The guard runs BEFORE the google-cloud-run import
(mirrors the
194         project guard), so it is exercisable in CI (the SDK is intentionally not a test
dependency)."""
195         from backend.compute.cloud_run import GoogleCloudRunJobClient
196
197         client = GoogleCloudRunJobClient()
198         for bogus in ("", "execution", "exec-123"):
199             with pytest.raises(RuntimeError, match="fully-qualified"):
200                 client.cancel_execution(bogus)
201
202
203         # ----- CloudRunCompute dispatch
+ poll
204
205
206     def test_cloud_run_dispatch_polls_and_returns_result(tmp_path):
207         client = _FakeRunClient(video_id="abc123", segment_count=3)
208         backend = CloudRunCompute(
209             run_client=client,
210             job_name="rally-worker",
211             region="asia-south1",
212             project="proj",
213             storage_config={},
214             policy=_FAST,
215             token="tok",
216         )
217         out = str(tmp_path / "bucket")
218         result = backend.run_infer(
219             ComputeJobSpec(
220                 video_uri="gs://b/videos/clip.mp4", out_uri=out, segmenter="wasb_hybrid"
221             )
222         )
223
224         assert result.returncode == 0
225         assert result.video_id == "abc123"
226         assert result.outputs_uri == f"{out}/outputs/abc123/"
227         assert result.report.get("segment_count") == 3
228         # exactly one Cloud Run Job execution was triggered, in the right region/project
229         assert len(client.calls) == 1
230         assert (
231             client.calls[0]["job"] == "rally-worker"
232             and client.calls[0]["region"] == "asia-south1"
233         )
234
235
236     def test_dispatch_argv_forces_inline_and_carries_overrides(tmp_path):
237         """The dispatched container must run INLINE (never re-dispatch) ?
compute_target=local_gpu +
238         native detector appended AFTER the user overrides (so they win); --done-uri always
present."""
239         client = _FakeRunClient()
240         backend = CloudRunCompute(
241             run_client=client, job_name="j", region="r", policy=_FAST, token="tok"
242         )
243         backend.run_infer(
244             ComputeJobSpec(
245                 video_uri="gs://b/v.mp4",
246                 out_uri=str(tmp_path / "bk"),
247                 segmenter="wasb_hybrid",
248                 config_overrides=("sport=tennis",),
249             )

```

```

250     )
251
252     argv = client.calls[0]["argv"]
253     assert argv[0] == "infer"
254     assert "--done-uri" in argv and "--video-uri" in argv and "--out-uri" in argv
255     assert argv[argv.index("--segmenter") + 1] == "wasb_hybrid"
256     # A DEFAULT spec (skip_validation=False) must NOT carry --skip-validation: the
worker still
257     # validates unless the dispatcher explicitly opts out. Guards the `if
spec.skip_validation:` gate
258     # against a refactor that appends the flag unconditionally (which would silently
bypass the direct
259     # `worker infer` cloud path's validation ? the very thing #skip-validation is scoped
away from).
260     assert "--skip-validation" not in argv
261     # the container-inline overrides come last (win) and the user override is carried
through
262     joined = " ".join(argv)
263     assert "deployment.compute_target=local_gpu" in joined
264     assert "indexing.detector_impl=native" in joined
265     assert "sport=tennis" in joined
266     # user override precedes the inline-forcing overrides in --set order
267     assert joined.index("sport=tennis") < joined.index(
268         "deployment.compute_target=local_gpu"
269     )
270
271
272     def test_dispatch_argv_carries_skip_validation_when_set(tmp_path):
273         """The builder must emit --skip-validation when the spec opts out ? the control-
plane sets
274         skip_validation=True on every dispatch (it is the authoritative validation gate),
and the worker
275         honours the flag to skip its redundant re-validation. Positive builder-level guard,
independent of
276         the control-plane path (which hardcodes True)."""
277         client = _FakeRunClient()
278         backend = CloudRunCompute(
279             run_client=client, job_name="j", region="r", policy=_FAST, token="tok"
280         )
281         backend.run_infer(
282             ComputeJobSpec(
283                 video_uri="gs://b/v.mp4",
284                 out_uri=str(tmp_path / "bk"),
285                 skip_validation=True,
286             )
287         )
288         assert "--skip-validation" in client.calls[0]["argv"]
289
290
291     def test_unique_token_per_call_when_not_injected(tmp_path):
292         """No injected token ? a fresh uuid per dispatch, so concurrent jobs don't collide
on the marker."""
293         client = _FakeRunClient()
294         backend = CloudRunCompute(run_client=client, job_name="j", region="r", policy=_FAST)
295         backend.run_infer(
296             ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bk"))
297         )
298         backend.run_infer(
299             ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bk"))
300         )
301         done1 = client.calls[0]["argv"][client.calls[0]["argv"].index("--done-uri") + 1]
302         done2 = client.calls[1]["argv"][client.calls[1]["argv"].index("--done-uri") + 1]
303         assert done1 != done2
304
305

```

```

306     def test_read_json_passes_dst_for_cloud_backends(monkeypatch):
307         """BLOCKER guard: GCS/S3 fetch_to_local REQUIRES a dst (a None dst raises) ?
_read_json must
308         pass one, else every real cloud marker/report read silently returns {} (local hid
this via its
309         identity passthrough). Mimic a cloud backend that rejects dst=None."""
310         from backend.compute import cloud_run as cr
311
312         captured = {}
313
314         def fake_fetch(uri, dst=None, *, config=None, overwrite=False):
315             captured["dst"] = dst
316             if dst is None:
317                 raise ValueError(
318                     "cloud backend requires a dst path"
319                 ) # mirror GCSStorage/S3Storage
320             with open(dst, "w", encoding="utf-8") as f:
321                 f.write({'video_id': "v", "returncode": 0})
322             return dst
323
324         monkeypatch.setattr(cr, "fetch", fake_fetch)
325         backend = cr.CloudRunCompute(run_client=_FakeRunClient(), job_name="j", region="r")
326         data = backend._read_json("gs://b/x.json")
327         assert (
328             captured["dst"] is not None
329         ) # the fix: a real dst is always passed (no silent {})
330         assert data == {"video_id": "v", "returncode": 0}
331         assert not os.path.exists(os.path.dirname(captured["dst"]))
332
333
334     def test_cloud_run_polls_until_marker_appears(monkeypatch):
335         """The bucket-poll WAIT path (D7 scale-to-zero cold start): the marker is ABSENT for
the first
336         polls then appears ? poll_until must LOOP, not resolve on poll 0. The synchronous
fake elsewhere
337         never exercises this; here exists() is False twice then True."""
338         from backend.compute import cloud_run as cr
339
340         exists_calls = {"n": 0}
341
342         class _FakeStore:
343             def exists(self, ref):
344                 exists_calls["n"] += 1
345                 return exists_calls["n"] >= 3 # absent for polls 1-2, present on poll 3
346
347         monkeypatch.setattr(cr, "get_storage", lambda backend, config=None: _FakeStore())
348
349         def fake_fetch(uri, dst=None, *, config=None, overwrite=False):
350             if dst is None:
351                 dst = os.path.join(tempfile.mkdtemp(prefix="cr-poll-"), "m.json")
352             with open(dst, "w", encoding="utf-8") as f:
353                 f.write({'video_id': "vlate", "returncode": 0, "segment_count": 1})
354             return dst
355
356         monkeypatch.setattr(cr, "fetch", fake_fetch)
357         backend = cr.CloudRunCompute(
358             run_client=_NoopRunClient(),
359             job_name="j",
360             region="r",
361             token="t",
362             policy=ExecutionPolicy(
363                 poll_every_n_sec=0.0, op_timeout_sec=5.0, max_wait_sec=5.0
364             ),
365         )
366         result = backend.run_infer(

```



```

367         ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri="gs://b")
368     )
369     assert result.video_id == "vlate"
370     assert (
371         exists_calls["n"] >= 3
372     ) # >=1 absent poll occurred before the marker appeared (loop ran)
373
374
375     # ----- poll timeout: rescue-or-cancel (R1-21 GPU
cost leak)
376
377     # max_wait_sec=0.0 ? poll_until gives up after its FIRST absent check (no real waiting).
378     _TIMEOUT_NOW = ExecutionPolicy(poll_every_n_sec=0.0, op_timeout_sec=5.0,
max_wait_sec=0.0)
379
380
381     def test_poll_timeout_cancels_execution_and_raises(tmp_path):
382         """Poll deadline + no marker ? the still-running execution is cancelled (with the
EXACT name
383         run_job returned) and PolicyTimeout surfaces carrying that name for the operator."""
384         client = _NeverDoneClient()
385         backend = CloudRunCompute(
386             run_client=client,
387             job_name="rally-worker",
388             region="asia-south1",
389             policy=_TIMEOUT_NOW,
390             token="tok",
391         )
392         with pytest.raises(PolicyTimeout) as exc:
393             backend.run_infer(
394                 ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bucket"))
395             )
396         assert client.cancelled == [_NeverDoneClient.EXEC_NAME] # cancel got the right name
397         msg = str(exc.value)
398         assert _NeverDoneClient.EXEC_NAME in msg # operator can find the execution in GCP
399         assert "cancel requested" in msg
400         assert "max_wait_sec" in msg # the original timeout text is preserved, not masked
401
402
403     def test_poll_timeout_cancel_failure_still_raises_policy_timeout(tmp_path):
404         """Best-effort means BEST-EFFORT: a cancel API failure must never mask the original
timeout ?
405         the caller (main.py's dispatch handler) still sees PolicyTimeout, with the failure
noted."""
406         client = _NeverDoneClient(cancel_error=RuntimeError("cancel API unreachable"))
407         backend = CloudRunCompute(
408             run_client=client,
409             job_name="rally-worker",
410             region="asia-south1",
411             policy=_TIMEOUT_NOW,
412             token="tok",
413         )
414         with pytest.raises(PolicyTimeout) as exc:
415             backend.run_infer(
416                 ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bucket"))
417             )
418         assert client.cancelled == [_NeverDoneClient.EXEC_NAME] # the attempt WAS made
419         msg = str(exc.value)
420         assert _NeverDoneClient.EXEC_NAME in msg
421         assert "cancel API unreachable" in msg # the cancel failure is surfaced, not
swallowed
422         assert "max_wait_sec" in msg # ...alongside the original timeout text
423
424
425     def test_poll_timeout_late_marker_rescues_result(monkeypatch):

```

```

426         """A slow-but-healthy worker that finishes JUST past the poll deadline
(max_wait_sec=1800 can
427         undercut the job's --task-timeout=3600) is rescued by the ONE final post-timeout
marker check:
428         the completed GPU result is ingested, nothing is cancelled, no timeout surfaces."""
429         from backend.compute import cloud_run as cr
430
431         exists_calls = {"n": 0}
432
433         class _FakeStore:
434             def exists(self, ref):
435                 exists_calls["n"] += 1
436                 return exists_calls["n"] >= 2 # absent for the poll, present on the final
check
437
438         monkeypatch.setattr(cr, "get_storage", lambda backend, config=None: _FakeStore())
439
440         def fake_fetch(uri, dst=None, *, config=None, overwrite=False):
441             if dst is None:
442                 dst = os.path.join(tempfile.mkdtemp(prefix="cr-late-"), "m.json")
443                 with open(dst, "w", encoding="utf-8") as f:
444                     f.write('{"video_id": "vlate", "returncode": 0, "segment_count": 1}')
445                 return dst
446
447         monkeypatch.setattr(cr, "fetch", fake_fetch)
448         client = _NoopRunClient()
449         backend = cr.CloudRunCompute(
450             run_client=client, job_name="j", region="r", token="t", policy=_TIMEOUT_NOW
451         )
452         result = backend.run_infer(
453             ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri="gs://b")
454         )
455         assert result.returncode == 0
456         assert result.video_id == "vlate" # the late result is ingested, not wasted
457         assert client.cancelled == [] # a finished execution is never cancelled
458         assert exists_calls["n"] == 2 # exactly ONE extra post-timeout check
459
460
461         # ----- execution-status fast negative path (worker crashed
pre-marker)
462
463         # poll_every_n_sec=0.0 ? the poll spins as fast as possible; max_wait_sec=2.0 BOUNDS any
regression
464         # (a fast path that silently fell back to marker-only polling fails in ~2s here, never a
65-min hang)
465         # while the working fast path aborts on poll 0 ? instantly.
466         _FAILFAST = ExecutionPolicy(poll_every_n_sec=0.0, op_timeout_sec=5.0, max_wait_sec=2.0)
467
468
469         def test_run_infer_fails_fast_when_execution_terminal_without_marker(tmp_path):
470             """The core fix: a worker execution that TERMINATED (Failed) without ever writing a
marker is
471             caught by the Executions-API probe ? RemoteExecutionFailed FAST (poll 0), not a
~65-min opaque
472             hang. The diagnostic names the execution, the terminal state, and the log-inspection
command."""
473             client = _CrashedNoMarkerClient(state="Failed")
474             backend = CloudRunCompute(
475                 run_client=client,
476                 job_name="rally-worker",
477                 region="asia-south1",
478                 policy=_FAILFAST,
479                 token="tok",
480             )
481             with pytest.raises(RemoteExecutionFailed) as exc:

```

```

482         backend.run_infer(
483             ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bucket"))
484         )
485     msg = str(exc.value)
486     assert _CrashedNoMarkerClient.EXEC_NAME in msg # operator can find the execution
487     assert "Failed" in msg # the terminal state is surfaced
488     assert "gcloud run jobs executions describe" in msg # actionable log pointer
489     assert client.status_polls == 1 # aborted on the FIRST terminal observation (fast,
not spun)
490     assert client.cancelled == [] # a finished execution is never cancelled (nothing to
cancel)
491
492
493     def test_dispatch_remote_maps_terminal_without_marker_to_rc5(tmp_path):
494         """The dispatch handler maps RemoteExecutionFailed to a distinct, clearly-labelled
rc 5 ? NOT
495         the rc 4 poll-timeout path (a still-maybe-running execution) and NOT a raw traceback
? so
496         'worker crashed before the marker' is operationally distinguishable and surfaces
immediately."""
497         import argparse
498
499         from backend.worker import __main__ as wmain
500
501         client = _CrashedNoMarkerClient(state="Failed")
502         backend = CloudRunCompute(
503             run_client=client,
504             job_name="rally-worker",
505             region="asia-south1",
506             policy=_FAILFAST,
507             token="tok",
508         )
509         args = argparse.Namespace(
510             video_uri="gs://b/v.mp4",
511             out_uri=str(tmp_path / "bucket"),
512             segmenter="wasb_hybrid",
513             set=None,
514         )
515         rc = wmain._dispatch_remote(backend, args)
516         assert rc == 5 # the crashed-pre-marker fast-fail code...
517         assert rc != 4 # ...distinct from the poll-budget timeout
518         assert client.status_polls == 1 # reached the fast path (not the timeout branch)
519
520
521     def test_succeeded_state_without_marker_also_fails_fast(tmp_path):
522         """Generalization: ANY terminal state with no marker is a guaranteed-futile wait. A
'Succeeded'
523         execution that nonetheless wrote no marker (the marker upload itself failed) also
fails fast,
524         with the state surfaced ? the whole 'execution over, no marker' class collapses, not
just Failed."""
525         client = _CrashedNoMarkerClient(state="Succeeded")
526         backend = CloudRunCompute(
527             run_client=client,
528             job_name="rally-worker",
529             region="asia-south1",
530             policy=_FAILFAST,
531             token="tok",
532         )
533         with pytest.raises(RemoteExecutionFailed) as exc:
534             backend.run_infer(
535                 ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bucket"))
536             )
537         assert "Succeeded" in str(exc.value)
538

```

```

539
540     def test_present_marker_wins_over_terminal_status(tmp_path):
541         """The done-marker stays AUTHORITATIVE: a worker that wrote a marker ? even a
FAILURE marker
542         (rc 2) ? resolves through the normal path with its REAL returncode, never a
synthesized
543         RemoteExecutionFailed, even when the execution status also reads terminal. Status is
a negative
544         path only; it must never override a real marker (which is checked first)."""
545
546         class _MarkerAndFailedStatus(_FakeRunClient):
547             def execution_terminal_state(self, execution): # pragma: no cover - never
reached
548                 return "Failed" # the marker resolves poll 0, so this is never consulted
549
550         client = _MarkerAndFailedStatus(video_id="abc", returncode=2)
551         backend = CloudRunCompute(
552             run_client=client, job_name="j", region="r", policy=_FAST, token="tok"
553         )
554         result = backend.run_infer(
555             ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bucket"))
556         )
557         assert result.returncode == 2 # the worker's REAL rc, not RemoteExecutionFailed
558         assert result.video_id == "abc"
559
560
561     def test_status_probe_error_is_swallowed_and_falls_back_to_marker_polling(tmp_path):
562         """Best-effort probe: a raising Executions API must NEVER fail a healthy run. The
wait degrades
563         to today's marker-only polling ? PolicyTimeout (the existing timeout path), NOT
564         RemoteExecutionFailed ? a flaky status API cannot manufacture a failure."""
565
566         class _RaisingStatusClient(_NeverDoneClient):
567             def execution_terminal_state(self, execution):
568                 raise RuntimeError("executions API unreachable")
569
570         client = _RaisingStatusClient()
571         backend = CloudRunCompute(
572             run_client=client,
573             job_name="rally-worker",
574             region="asia-south1",
575             policy=_TIMEOUT_NOW,
576             token="tok",
577         )
578         with pytest.raises(PolicyTimeout): # swallowed probe error ? the unchanged timeout
path
579             backend.run_infer(
580                 ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bucket"))
581             )
582         assert client.cancelled == [_NeverDoneClient.EXEC_NAME] # still cancels on the real
timeout
583
584
585     def test_default_execution_terminal_state_is_none():
586         """The ABC default is None (additive / default-OFF): a client that doesn't override
status
587         polling degrades transparently to marker-only behaviour, so every pre-existing fake
is
588         unaffected and the seam never changes a healthy run's outcome."""
589         assert _NoopRunClient().execution_terminal_state("whatever") is None
590         assert (
591             _NeverDoneClient().execution_terminal_state(_NeverDoneClient.EXEC_NAME) is None
592         )
593
594

```

```

595     def test_real_client_terminal_state_none_for_unresolvable_name():
596         """The real client's status probe is best-effort: a non-resolvable execution name
returns None
597         (can't probe ? keep marker-polling), NOT a raise ? exercisable in CI (the SDK is
absent)."""
598         from backend.compute.cloud_run import GoogleCloudRunJobClient
599
600         client = GoogleCloudRunJobClient()
601         for bogus in ("", "execution", "exec-123"):
602             assert client.execution_terminal_state(bogus) is None
603
604
605     def test_terminal_status_with_marker_written_in_race_resolves_via_marker(monkeypatch):
606         """Concurrency race: the execution reads terminal exactly as the worker finishes
writing its
607         marker. _poll_check re-checks the marker ONCE after a terminal status ? if it landed
in between,
608         the completed result is ingested (the marker wins), NOT a spurious
RemoteExecutionFailed."""
609         from backend.compute import cloud_run as cr
610
611         exists_calls = {"n": 0}
612
613         class _FakeStore:
614             def exists(self, ref):
615                 exists_calls["n"] += 1
616                 # absent on the pre-status read, present on the ONE post-status race re-
check
617                 return exists_calls["n"] >= 2
618
619         monkeypatch.setattr(cr, "get_storage", lambda backend, config=None: _FakeStore())
620
621         def fake_fetch(uri, dst=None, *, config=None, overwrite=False):
622             if dst is None:
623                 dst = os.path.join(tempfile.mkdtemp(prefix="cr-race-"), "m.json")
624                 with open(dst, "w", encoding="utf-8") as f:
625                     f.write('{"video_id": "vrace", "returncode": 0, "segment_count": 1}')
626                 return dst
627
628         monkeypatch.setattr(cr, "fetch", fake_fetch)
629         client = _CrashedNoMarkerClient(state="Failed") # status says terminal-Failed...
630         backend = cr.CloudRunCompute(
631             run_client=client, job_name="j", region="r", token="t", policy=_FAILFAST
632         )
633         result = backend.run_infer(
634             ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri="gs://b")
635         )
636         assert result.video_id == "vrace" # ...but the marker that raced in WINS
637         assert exists_calls["n"] == 2 # exactly the pre-status read + the ONE race re-check
638         assert client.status_polls == 1 # status was consulted once, then the marker
superseded it
639
640
641     def test_classify_execution_state_maps_proto_plus_fields():
642         """Regression guard for the proto-plus `.seconds` bug: google-cloud-run marshals a
Timestamp to a
643         datetime (which has `.second`, NOT `.seconds`) when set and to None when unset.
Terminality MUST
644         key on `completion_time is None`, and the int per-task counts classify it. The
original
645         `getattr(completion_time, "seconds", 0)` guard was ALWAYS 0 ? the real fast-fail
path was inert."""
646         from datetime import datetime, timezone
647
648         from backend.compute.cloud_run import _classify_execution_state

```

```

649
650     class _Exe: # mimics a proto-plus run_v2.Execution (datetime completion_time + int
counts)
651         def __init__(
652             self, completion_time=None, cancelled_count=0, failed_count=0,
succeeded_count=0
653         ):
654             self.completion_time = completion_time
655             self.cancelled_count = cancelled_count
656             self.failed_count = failed_count
657             self.succeeded_count = succeeded_count
658
659         done = datetime(2026, 7, 7, 12, 0, tzinfo=timezone.utc)
660         # a datetime has `.second`, never `.seconds` ? the exact reason the old guard mis-
fired
661         assert hasattr(done, "second") and not hasattr(done, "seconds")
662         # still pending/running (no completion timestamp) ? None, even with a stray count
set
663         assert _classify_execution_state(_Exe(completion_time=None)) is None
664         assert _classify_execution_state(_Exe(completion_time=None, failed_count=1)) is None
665         # terminal classifications
666         assert _classify_execution_state(_Exe(completion_time=done, failed_count=1)) ==
"Failed"
667         assert (
668             _classify_execution_state(_Exe(completion_time=done, cancelled_count=1)) ==
"Cancelled"
669         )
670         assert (
671             _classify_execution_state(_Exe(completion_time=done, succeeded_count=1)) ==
"Succeeded"
672         )
673         # cancel takes precedence over a co-incident failed_count
674         assert (
675             _classify_execution_state(
676                 _Exe(completion_time=done, cancelled_count=1, failed_count=1)
677             )
678             == "Cancelled"
679         )
680         # terminal but no positive count ? conservative Failed
681         assert _classify_execution_state(_Exe(completion_time=done)) == "Failed"
682
683
684     def test_classify_execution_state_against_real_proto_plus():
685         """Belt-and-suspenders: exercise the classifier against a REAL google-cloud-run
run_v2.Execution
686         so a future proto-plus marshaling change is caught. Skipped where the SDK isn't
installed (it's
687         intentionally NOT a CI dependency); the fake-based test above pins the shape for
CI."""
688         run_v2 = pytest.importorskip("google.cloud.run_v2")
689         from datetime import datetime, timezone
690
691         from backend.compute.cloud_run import _classify_execution_state
692
693         ts = datetime(2026, 7, 7, 12, 0, tzinfo=timezone.utc)
694         running = run_v2.Execution() # completion_time unset ? proto-plus marshals to None
695         assert _classify_execution_state(running) is None
696
697         for count_attr, expected in (
698             ("failed_count", "Failed"),
699             ("cancelled_count", "Cancelled"),
700             ("succeeded_count", "Succeeded"),
701         ):
702             exe = run_v2.Execution()
703             exe.completion_time = ts

```

```

704         setattr(exe, count_attr, 1)
705         assert _classify_execution_state(exe) == expected
706
707
708     def test_terminal_execution_state_none_for_empty_execution():
709         """An empty/unknown execution name (run_job could not resolve one) can't be probed ?
None
710         (degrade to marker-only polling), and the client's status API is never even
called."""
711
712         class _Boom(_NoopRunClient):
713             def execution_terminal_state(self, execution): # pragma: no cover - must never
be called
714                 raise AssertionError("must not probe an empty execution name")
715
716         backend = CloudRunCompute(run_client=_Boom(), job_name="j", region="r")
717         assert backend._terminal_execution_state("") is None
718
719
720     # ----- worker wiring: dispatch branch +
done-marker
721
722
723     def test_cmd_infer_dispatches_when_cloud_target(tmp_path, monkeypatch):
724         """cmd_infer with compute_target=cloud_run hands off to the ComputeBackend (no in-
process run)."""
725         from backend.compute.base import ComputeResult
726         from backend.worker import __main__ as wmain
727
728         seen = {}
729
730         class _FakeBackend:
731             def run_infer(self, spec):
732                 seen["spec"] = spec
733                 return ComputeResult(
734                     returncode=0, video_id="vidX", outputs_uri="gs://b/outputs/vidX/"
735                 )
736
737         monkeypatch.setattr(wmain, "get_compute_backend", lambda config: _FakeBackend())
738         # cloud-serving startup checks require the AI-handoff key (AI_HANDOFF_KEY_MISSING is
fatal).
739         monkeypatch.setenv("GEMINI_API_KEY", "test-key")
740         cfg = tmp_path / "c.json"
741         cfg.write_text(
742             '{"deployment": {"profile": "cloud-serving"}, "storage": {"backend": "gcs"}}'
743         )
744         rc = wmain.cmd_infer(
745             wmain.build_parser().parse_args(
746                 [
747                     "infer",
748                     "--video-uri",
749                     "gs://b/videos/clip.mp4",
750                     "--out-uri",
751                     "gs://b",
752                     "--config",
753                     str(cfg),
754                     "--segmenter",
755                     "wasb_hybrid",
756                 ]
757             )
758         )
759         assert rc == 0
760         assert seen["spec"].video_uri == "gs://b/videos/clip.mp4"
761         assert seen["spec"].segmenter == "wasb_hybrid"
762

```

```

763
764     def test_worker_writes_done_marker_on_success(tmp_path, monkeypatch):
765         """The container side: cmd_infer with --done-uri writes a completion marker
766         (video_id + rc) so a
767         remote dispatcher's bucket-poll terminates. Driven by the deterministic stub + skip-
768         AI."""
769         from backend.worker import __main__ as wmain
770
771         class _OK:
772             passed = True
773             validator_name = "fake"
774             message = "ok"
775
776             def model_dump(self):
777                 return {"validator_name": "fake", "passed": True, "message": "ok"}
778
779         monkeypatch.setattr(
780             wmain.validator_registry,
781             "run_all_with_context",
782             lambda path, cfg: ([_OK()], {}),
783         )
784         monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
785         monkeypatch.setattr(
786             "backend.pipeline.segmenters.trajectory_hybrid.get_video_duration",
787             lambda p: 30.0,
788         )
789
790         cfg = tmp_path / "c.json"
791         cfg.write_text(
792             json.dumps(
793                 {
794                     "indexing": {"detector_impl": "stub", "skip_ai_handoff": True},
795                     "output": {"output_dir": str(tmp_path / "o")},
796                 }
797             )
798         )
799         video = tmp_path / "in.mp4"
800         video.write_bytes(b"FAKEVIDEO")
801         out = tmp_path / "bucket"
802         done = str(out / "_dispatch" / "tok.done")
803
804         rc = wmain.cmd_infer(
805             wmain.build_parser().parse_args(
806                 [
807                     "infer",
808                     "--video-uri",
809                     str(video),
810                     "--out-uri",
811                     str(out),
812                     "--config",
813                     str(cfg),
814                     "--scratch",
815                     str(tmp_path / "s"),
816                     "--segmenter",
817                     "wasb_hybrid",
818                     "--done-uri",
819                     done,
820                 ]
821             )
822         )
823         assert rc == 0
824         marker = json.loads(open(done, encoding="utf-8").read())
825         assert marker["returncode"] == 0 and marker["status"] == "success"
826         assert marker["video_id"] and marker["segment_count"] >= 1
827         # CP-rebuild P3 (#485): the report records the SOURCE the detection ran on, so a

```



```

826     # restart-wiped control plane re-hydrating from outputs/<md5>/ can restore the video
827     # row's truthful path (not just the segments).
828     report_path = out / "outputs" / marker["video_id"] / "report.json"
829     report = json.loads(report_path.read_text(encoding="utf-8"))
830     assert report["video_uri"] == str(video)
831
832
833     def test_worker_no_done_uri_is_unchanged(tmp_path, monkeypatch):
834         """DEFAULT-OFF: without --done-uri the inline run writes NO marker (today's
behaviour)."""
835         from backend.worker import __main__ as wmain
836
837         class _OK:
838             passed = True
839             validator_name = "fake"
840             message = "ok"
841
842             def model_dump(self):
843                 return {"validator_name": "fake", "passed": True, "message": "ok"}
844
845         monkeypatch.setattr(
846             wmain.validator_registry,
847             "run_all_with_context",
848             lambda path, cfg: ([_OK()], {}),
849         )
850         monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
851         monkeypatch.setattr(
852             "backend.pipeline.segmenters.trajectory_hybrid.get_video_duration",
853             lambda p: 30.0,
854         )
855         cfg = tmp_path / "c.json"
856         cfg.write_text(
857             json.dumps(
858                 {
859                     "indexing": {"detector_impl": "stub", "skip_ai_handoff": True},
860                     "output": {"output_dir": str(tmp_path / "o")},
861                 }
862             )
863         )
864         video = tmp_path / "in.mp4"
865         video.write_bytes(b"FAKEVIDEO")
866         out = tmp_path / "bucket"
867         rc = wmain.cmd_infer(
868             wmain.build_parser().parse_args(
869                 [
870                     "infer",
871                     "--video-uri",
872                     str(video),
873                     "--out-uri",
874                     str(out),
875                     "--config",
876                     str(cfg),
877                     "--scratch",
878                     str(tmp_path / "s"),
879                     "--segmenter",
880                     "wasb_hybrid",
881                 ]
882             )
883         )
884         assert rc == 0
885         assert not (out / "_dispatch").exists() # no marker dir when --done-uri is omitted
886
887
888     def test_worker_writes_failure_marker_on_validation_fail(tmp_path, monkeypatch):
889         """The container side: a worker-side FAILURE (validation rc=2) with --done-uri

```

```

writes a marker
890     carrying the REAL rc, so a remote dispatcher's bucket-poll terminates FAST (not a
30-min hang)."""
891     from backend.worker import __main__ as wmain
892
893     class _BAD:
894         passed = False
895         validator_name = "blur"
896         message = "blurry"
897
898         def model_dump(self):
899             return {"validator_name": "blur", "passed": False, "message": "blurry"}
900
901     monkeypatch.setattr(
902         wmain.validator_registry,
903         "run_all_with_context",
904         lambda path, cfg: ([_BAD()], {}),
905     )
906     monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
907     cfg = tmp_path / "c.json"
908     cfg.write_text(
909         json.dumps(
910             {
911                 "indexing": {"detector_impl": "stub", "skip_ai_handoff": True},
912                 "output": {"output_dir": str(tmp_path / "o")},
913             }
914         )
915     )
916     video = tmp_path / "in.mp4"
917     video.write_bytes(b"FAKEVIDEO")
918     out = tmp_path / "bucket"
919     done = str(out / "_dispatch" / "tok.done")
920
921     rc = wmain.cmd_infer(
922         wmain.build_parser().parse_args(
923             [
924                 "infer",
925                 "--video-uri",
926                 str(video),
927                 "--out-uri",
928                 str(out),
929                 "--config",
930                 str(cfg),
931                 "--scratch",
932                 str(tmp_path / "s"),
933                 "--segmenter",
934                 "wasb_hybrid",
935                 "--done-uri",
936                 done,
937             ]
938         )
939     )
940     assert rc == 2 # validation failure
941     marker = json.loads(open(done, encoding="utf-8").read())
942     assert marker["returncode"] == 2 and marker["status"] == "failed"
943     assert marker[
944         "video_id"
945     ] # carries the id so the dispatcher's poll terminates (rc != 0 ? failed)
946
947
948     def test_worker_failure_writes_no_marker_without_done_uri(tmp_path, monkeypatch):
949         """DEFAULT-OFF: a worker failure with NO --done-uri writes no marker (the unchanged
local run)."""
950         from backend.worker import __main__ as wmain
951

```

```

952 class _BAD:
953     passed = False
954     validator_name = "blur"
955     message = "blurry"
956
957     def model_dump(self):
958         return {"validator_name": "blur", "passed": False, "message": "blurry"}
959
960 monkeypatch.setattr(
961     wmain.validator_registry,
962     "run_all_with_context",
963     lambda path, cfg: ([_BAD()], {}),
964 )
965 monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
966 cfg = tmp_path / "c.json"
967 cfg.write_text(
968     json.dumps(
969         {
970             "indexing": {"detector_impl": "stub", "skip_ai_handoff": True},
971             "output": {"output_dir": str(tmp_path / "o")},
972         }
973     )
974 )
975 video = tmp_path / "in.mp4"
976 video.write_bytes(b"FAKEVIDEO")
977 out = tmp_path / "bucket"
978 rc = wmain.cmd_infer(
979     wmain.build_parser().parse_args(
980         [
981             "infer",
982             "--video-uri",
983             str(video),
984             "--out-uri",
985             str(out),
986             "--config",
987             str(cfg),
988             "--scratch",
989             str(tmp_path / "s"),
990             "--segmenter",
991             "wasb_hybrid",
992         ]
993     )
994 )
995 assert rc == 2
996 assert not (out / "_dispatch").exists()
997
998
999 # ----- Dockerfile structural guard
1000 (no build)
1001
1002 def test_dockerfile_has_required_directives():
1003     """CI can't build a CUDA image, so guard the Dockerfile's load-bearing directives
1004     structurally:
1005     the GPU base, ffmpeg + DejaVu fonts (the watermark), the WASB conda env, and the
1006     worker entrypoint."""
1007     from pathlib import Path
1008
1009     text = Path("deploy/cloudrun/Dockerfile").read_text(encoding="utf-8")
1010     assert "FROM nvidia/cuda" in text # GPU base (CUDA)
1011     assert "ffmpeg" in text # stitch/decode
1012     assert "fonts-dejavu" in text # the watermark filter font
1013     # The #506 modernized WASB stack (L4-validated): torch 2.6 cu124 + PyNvVideoCodec,
1014     and
1015     # the NVDEC/NVENC driver userspace libs PyNvVideoCodec dlopens (compute-only GPU

```

```

hosts
1013     # inject neither) ? a silent drop of any of these bricks
decode_backend=nvdec_resident.
1014     assert "torch==2.6.0" in text
1015     assert "download.pytorch.org/whl/cu124" in text
1016     assert "PyNvVideoCodec" in text
1017     assert "libnvcuvid.so.1" in text and "libnvidia-encode.so.1" in text
1018     assert "torch==1.11.0+cu113" not in text # the old stack is gone, not co-installed
1019     assert "WASB-SBDT" in text # the shuttle detector repo
1020     assert 'ENTRYPOINT ["python3", "-m", "backend.worker"]' in text
1021     assert "WASB_PYTHON=/opt/conda/envs/wasb/bin/python" in text
1022     # The cloud worker reads the gs:// input + writes the gs:// outputs/done-marker
through
1023     # backend.storage (storage-gcs); and when this image is reused as the deployed
control plane it
1024     # dispatches the L4 Job via google.cloud.run_v2 (cloud-run) and verifies Cf-Access
JWTs when
1025     # edge-auth is on (auth ? PyJWT). A bare `pip install .` boots but dies on the first
gs:// fetch
1026     # (ModuleNotFoundError: google.cloud), on dispatch (google-cloud-run missing), or on
the first
1027     # edge-authed request (jwt missing) ? real Cloud Run failures the mock-client tests
can't see.
1028     # Lock the extras in structurally.
1029     assert ".[storage-gcs,cloud-run,auth]" in text
1030
1031
1032     def test_dockerfile_uses_micromamba_not_anaconda_tos():
1033         """P3a: the WASB env is built with micromamba + conda-forge (BSD, no ToS), NOT
Miniconda +
1034         `conda tos accept` (the ?200-employee Anaconda Terms-of-Service trap,
PACKAGING_PLAN.md §3).
1035         Locked structurally so a revert to the ToS path fails CI; the env path (WASB_PYTHON)
is unchanged."""
1036         from pathlib import Path
1037
1038         text = Path("deploy/cloudrun/Dockerfile").read_text(encoding="utf-8")
1039         assert "micromamba" in text # the conda-forge env builder
1040         assert "conda-forge" in text # the BSD, ToS-free channel
1041         assert (
1042             "MAMBA_ROOT_PREFIX=/opt/conda" in text
1043         ) # keeps the env at /opt/conda/envs/wasb
1044         assert "Miniconda3" not in text # the Anaconda installer is gone
1045         assert "conda tos accept" not in text # no ToS acceptance needed (conda-forge)
1046
1047
1048     def test_remote_poll_budget_comes_from_config_unless_policy_injected():
1049         """CP-rebuild §4.3: deployment.remote_poll_max_wait_sec drives the dispatch poll
budget
1050         (default 3900s > the Job's --task-timeout 3600s ? the old 1800s default could
abandon a
1051         slow-but-healthy paid GPU run), and an explicitly injected policy always wins."""
1052         from backend.infrastructure.execution_policy import DEFAULT_POLICY
1053
1054         cfg = AppConfig()
1055         cfg.deployment.compute_target = "cloud_run"
1056
1057         be = get_compute_backend(cfg, run_client=_FakeRunClient())
1058         assert be is not None
1059         assert be._policy.max_wait_sec == 3900.0
1060         # Only the budget is derived ? the other resilience knobs stay at their defaults.
1061         assert be._policy.poll_every_n_sec == DEFAULT_POLICY.poll_every_n_sec
1062         assert be._policy.op_timeout_sec == DEFAULT_POLICY.op_timeout_sec
1063
1064         cfg.deployment.remote_poll_max_wait_sec = 42.0

```

```

1065         assert get_compute_backend(cfg, run_client=_FakeRunClient())._policy.max_wait_sec ==
42.0
1066
1067         # An injected policy is never overridden by config.
1068         be = get_compute_backend(cfg, run_client=_FakeRunClient(), policy=_FAST)
1069         assert be._policy.max_wait_sec == _FAST.max_wait_sec
1070
1071
1072     def test_run_infer_forwards_on_wait_into_the_bucket_poll(tmp_path, monkeypatch):
1073         """The #482 heartbeat plumbing: the on_wait callback handed to run_infer must be the
1074         exact on_tick the bucket-poll drives (cloud_run threads it through poll_until)."""
1075         from backend.compute import cloud_run as cr
1076
1077         seen = {}
1078
1079         def fake_poll(check, policy, *, label="op", logger=None, on_tick=None, **kw):
1080             seen["on_tick"] = on_tick
1081             return {"video_id": "vidCloud", "returncode": 0}
1082
1083         monkeypatch.setattr(cr, "poll_until", fake_poll)
1084         client = _FakeRunClient() # still writes report.json so the post-marker read
succeeds
1085         backend = CloudRunCompute(
1086             run_client=client,
1087             job_name="rally-worker",
1088             region="asia-southeast1",
1089             storage_config={},
1090             policy=_FAST,
1091             token="tok",
1092         )
1093
1094         def heartbeat(_elapsed):
1095             pass
1096
1097         result = backend.run_infer(
1098             ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bk")),
1099             on_wait=heartbeat,
1100         )
1101         assert seen["on_tick"] is heartbeat
1102         assert result.video_id == "vidCloud"
1103

```

49. assistant (2026-07-08T05:17:02.444Z)

I have all the patterns I need. Let me check `backend/main.py` module-level side effects (to decide whether the worker can reuse the `\_main.VideoCompiler` seam) and how the compile job is tested.

50. user (2026-07-08T05:17:03.958Z)

```

1     import argparse
2     import json
3     import logging
4     import os
5     import shutil
6     import sys
7     import tempfile
8     import threading
9     import uuid
10    from importlib.metadata import PackageNotFoundError
11    from importlib.metadata import version as _pkg_version
12    from urllib.parse import quote
13
14    import uvicorn
15    from dotenv import load_dotenv

```

```

16
17     # PACKAGING_PLAN.md P0a: prefer an app-home .env ONLY when RALLY_APP_HOME is set (the
packaged
18     # launcher pins it); otherwise env_file_to_load() returns None and load_dotenv(None)
keeps the
19     # exact historical frame-relative walk-up ? byte-identical to the pre-P0a bare
load_dotenv().
20     # stdlib-only import ? keeps this pre-pydantic line light.
21     from backend.utils.app_home import env_file_to_load, resolve_output_dir
22
23     load_dotenv(env_file_to_load())
24
25     # Centralized logging setup for the indexer (hot paths use logger, CLI output uses print
for direct user-facing messages).
26     logging.basicConfig(
27         level=logging.INFO,
28         format="%(asctime)s [%(levelname)s] %(name)s: %(message)s",
29         datefmt="%H:%M:%S",
30     )
31     logger = logging.getLogger(__name__)
32     from typing import Any, Dict, List, Optional
33
34     from fastapi import FastAPI, HTTPException, Request
35     from fastapi.middleware.cors import CORSMiddleware
36     from fastapi.responses import (
37         FileResponse,
38         HTMLResponse,
39         JSONResponse,
40         RedirectResponse,
41     )
42     from fastapi.staticfiles import StaticFiles
43     from pydantic import BaseModel
44     from starlette.middleware.trustedhost import TrustedHostMiddleware
45
46     from backend import (
47         billing, # M8: per-job cost ledger (USD internal / INR display)
48         storage, # M2: URI-aware input acquisition (local = identity passthrough)
49     )
50     from backend.api import (
51         auth as edge_auth, # M9: Cloudflare Access edge-auth (default-OFF)
52     )
53     from backend.api import (
54         jobs as jobs_api, # /api/jobs list route (router) + submit-dedup helper (#448)
55     )
56     from backend.config import (
57         AppConfig,
58         StartupCheckError,
59         enforce_startup_checks,
60         write_light_default_config, # P2b: batteries-included first-run config seeding
61     )
62     from backend.config import ( # new typed config
63         load_config as load_app_config,
64     )
65     from backend.infrastructure.database import Database
66     from backend.pipeline.segmenters.base import segmenter_registry
67     from backend.pipeline.stitcher.compiler import VideoCompiler
68     from backend.pipeline.validators.base import registry
69     from backend.providers import provider_registry
70     from backend.reporting import emit_indexer_report
71     from backend.results import Failure # standardized error/result contract
72
73     # Server/network bootstrap helpers (CORS/host-allowlist config + the `indexer --app`
launcher)
74     # live in backend.server_runtime (P24 god-module split). Imported HERE, above `app =
FastAPI(...)`

```

```

75     # because the module-level add_middleware(...) below calls
_cors_origins_from_env/_allowed_hosts_from_env
76     # at import time; re-exported so `import backend.main as m; m._resolve_bind_host(...)`
keeps working.
77     from backend.server_runtime import ( # noqa: F401 (re-exports: names stay addressable
on backend.main)
78         _allowed_hosts_from_env,
79         _cors_origins_from_env,
80         _find_free_port,
81         _open_browser_when_ready,
82         _port_is_free,
83         _resolve_app_port,
84         _resolve_bind_host,
85     )
86     from backend.storage.capabilities import (
87         storage_capabilities, # P3: secret-free "which media can this box use" report
88     )
89     from backend.utils.ffmpeg import ( # P2a: single ffmpeg/ffprobe resolver
90         ffmpeg_bin,
91         ffprobe_bin,
92     )
93     from backend.utils.freespace import require_free_bytes # HTTP-507 disk guard (P2, D11)
94     from backend.utils.hashing import compute_video_id # canonical file-identity hashing
95     from backend.utils.paths import (
96         output_base,
97         resolve_under_root,
98         safe_output_filename,
99         validate_input_path,
100     )

```

51. user (2026-07-08T05:17:05.556Z)

```

480: def test_create_compile_job_persists_kind_and_params(tmp_path):
481:     """#329: an async compile job rides the jobs table with kind='compile' + a JSON params
payload
484:     assert db.create_compile_job(
488:         assert job["kind"] == "compile"
498: def test_worker_dispatches_compile_kind_to_execute_compile(env):
499:     """The claimed-job dispatcher routes kind='compile' to _execute_compile (NOT
_execute_processing),
512:     mp.setattr(m, "_execute_compile", _fake_compile)
518:     db.create_compile_job("cj2", "vidX", "/v.mp4", [1], "r.mp4")
567:     db.create_compile_job("cmp_q", "vid1", "/v.mp4", [1, 2], "reel_q.mp4")
568:     db.create_compile_job("cmp_r", "vid1", "/v.mp4", [1, 2], "reel_r.mp4")
570:     db.create_compile_job("cmp_done", "vid1", "/v.mp4", [1, 2], "reel_done.mp4")
683:     db.create_compile_job("c1", "vid", str(vid), [1], "reel.mp4") # active compile, owner
None

```

52. assistant (2026-07-08T05:17:33.537Z)

`backend.main` is importable without creating a DB (heavy but fine on the worker image). Let me check how `\_execute\_compile`'s stitch is currently tested (VideoCompiler seam) so my refactor keeps those green.

53. user (2026-07-08T05:17:36.420Z)

```

=== VideoCompiler / _execute_compile / _stitch usages in tests ===
tests/test_api_endpoints.py
tests/test_async_jobs.py
tests/test_compiler.py
tests/test_generate_highlights.py
tests/test_main_module_budget.py
tests/test_upload_download.py
tests/test_watermark.py

```

```

tests/test_watermark_i18n.py
tests/___pycache___/test_async_jobs.cpython-313-pytest-9.0.3.pyc
tests/___pycache___/test_compiler.cpython-313-pytest-9.0.3.pyc
=== compile docs dir ===
architecture.svg
README.md
runlogs
TESTING_STRATEGY.md
=== localize_watermark ===
"""Localized reel-watermark text + font selection.

```

The SPA passes the UI locale to ``POST /api/compile``; this maps it to a localized branding string and a script-appropriate bundled font, so a Kannada/Hindi/? user's reel carries the watermark in their language. The brand token ("KhelSutra.guru") stays Latin; only the surrounding phrase localizes.

Fonts (``backend/assets/fonts/``):

- en/es/da/id (Latin) ? the configured default (bundled RobotoSlab, Apache-2.0)
- kn ? Noto Sans Kannada · hi ? Noto Sans Devanagari (OFL-1.1, bundled ? the D7 OFL-for-media carve-out; FRIENDLY_AND_MULTILINGUAL §4 anticipated this).

Honest fallbacks: an unknown/missing locale (or a locale with no entry, incl. ``en``) returns the watermark UNCHANGED, so existing/default behavior is preserved. If a script font file is somehow absent, ``font\_path`` is left at the default (RobotoSlab) rather than pointing at a missing file ? and the compiler already retries the encode without the watermark if drawtext can't render, so a font gap can never nuke the reel. Kept config-package-agnostic (operates on a dict).

```

"""

```

```

from pathlib import Path
from typing import Any, Dict, Optional

```

This module is backend/pipeline/stitcher/ ? parents[2] == backend/ (matches compiler

```

_FONTS_DIR = Path(__file__).resolve().parents[2] / "assets" / "fonts"

```

Localized watermark phrase per locale. AI-drafted ? native-speaker review pending (sa the SPA catalogs). ``en`` is intentionally absent: it falls through to the configured default

```

WATERMARK_TEXT: Dict[str, str] = {
    "es": "Generado con KhelSutra.guru",
    "da": "Lavet med KhelSutra.guru",
    "id": "Dibuat dengan KhelSutra.guru",
    "kn": "KhelSutra.guru ??? ????",
    "hi": "KhelSutra.guru ?? ???? ??",
}

```

Script-specific bundled fonts ? only where the Latin default (RobotoSlab) can't render Latin locales (es/da/id) localize the TEXT but keep the default font.

```

FONT_FILE_BY_LOCALE: Dict[str, str] = {
    "kn": "NotoSansKannada-Variable.ttf",
    "hi": "NotoSansDevanagari-Variable.ttf",
}

```

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

54. user (2026-07-08T05:17:42.793Z)

```

1      """VideoCompiler coverage (previously zero ? the user-facing highlight-reel output path,
2      where the f-string SyntaxError lived undetected). ffmpeg/ffprobe are mocked."""
3
4      import os

```



```

5     import subprocess
6     from unittest.mock import MagicMock, patch
7
8     import pytest
9
10    from backend.pipeline.stitcher.compiler import VideoCompiler
11
12
13    def test_parse_time_formats():
14        f = VideoCompiler._parse_time
15        assert f(12.5) == 12.5
16        assert f("12.5") == 12.5
17        assert f("1:30") == 90.0
18        assert f("1:00:00") == 3600.0
19        assert f("garbage") == 0.0
20
21
22    def test_merge_overlapping_collapses_duplicates_and_overlaps():
23        m = VideoCompiler._merge_overlapping
24        # Exact duplicates + near-duplicates (the "every rally twice" re-run bug) collapse
to one each.
25        assert m([(28.2, 34.5), (28.5, 34.5), (40.5, 45.5), (40.5, 45.5)]) == [
26            (28.2, 34.5),
27            (40.5, 45.5),
28        ]
29        # Unsorted input is sorted; an overlapping range extends the previous.
30        assert m([(100.0, 110.0), (0.0, 5.0), (104.0, 120.0)]) == [
31            (0.0, 5.0),
32            (100.0, 120.0),
33        ]
34        # Genuinely distinct (gapped) rallies are preserved, not merged.
35        assert m([(0.0, 5.0), (6.0, 9.0)]) == [(0.0, 5.0), (6.0, 9.0)]
36        # Degenerate rows (end <= start) are dropped; string timestamps are accepted.
37        assert m([("0:10", "0:20"), (30.0, 30.0), (31.0, 29.0)]) == [(10.0, 20.0)]
38
39
40    def test_compile_dedups_before_stitching(tmp_path):
41        """A segment list with every rally twice must yield ONE clip per unique rally."""
42        comp = VideoCompiler(str(tmp_path / "out"))
43        raw = tmp_path / "raw.mp4"
44        raw.write_bytes(b"x")
45        trims = []
46
47        def fake_run(cmd, *a, **k):
48            if cmd[0] == "ffprobe":
49                return MagicMock(stdout="100.0\n", returncode=0)
50            out = cmd[-1]
51            if isinstance(out, str) and out.endswith(".mp4"):
52                open(out, "wb").write(b"clip")
53                # count per-clip trims (the concat list file input is not an .mp4 output)
54                if "-i" in cmd and cmd.index("-i") + 1 == str(raw):
55                    trims.append(cmd)
56            return MagicMock(returncode=0, stdout=b"", stderr=b"")
57
58        dup_segments = [(0.0, 5.0), (0.2, 5.0), (10.0, 15.0), (10.0, 15.0), (20.0, 25.0)]
59        with patch.object(subprocess, "run", side_effect=fake_run):
60            comp.compile_highlights(str(raw), dup_segments, "reel.mp4")
61        # 5 input segments collapse to 3 unique rallies -> 3 per-clip trims, not 5.
62        assert len(trims) == 3, f"expected 3 deduped trims, got {len(trims)}"
63
64
65    def test_compile_basenames_output_and_stitches(tmp_path):
66        comp = VideoCompiler(str(tmp_path / "out"))
67        raw = tmp_path / "raw.mp4"
68        raw.write_bytes(b"x")

```

```

69
70     calls = []
71
72     def fake_run(cmd, *a, **k):
73         calls.append(cmd)
74         # Create whatever output file ffmpeg/ffprobe was told to write so existence
checks pass.
75         if cmd[0] == "ffprobe":

```

55. user (2026-07-08T05:17:47.405Z)

```

495     assert db.get_job("pj")["kind"] is None
496
497
498     def test_worker_dispatches_compile_kind_to_execute_compile(env):
499         """The claimed-job dispatcher routes kind='compile' to _execute_compile (NOT
_execute_processing),
500         records the result as 'done', and skips the process-only cost/flywheel hooks."""
501         _client, db, _tmp_path, mp = env
502         seen = {}
503
504         def _fake_compile(cfg, db, job, progress=None):
505             seen["id"] = job["id"]
506             return {
507                 "status": "success",
508                 "video_id": "vidX",
509                 "download_url": "/api/highlights/vidX/r.mp4",
510             }
511
512         mp.setattr(m, "_execute_compile", _fake_compile)
513         mp.setattr(
514             m,
515             "_execute_processing",
516             lambda *a, **k: pytest.fail("a compile job must not hit _execute_processing"),
517         )
518         db.create_compile_job("cj2", "vidX", "/v.mp4", [1], "r.mp4")
519         m._run_claimed_job(db.get_job("cj2"), m.global_config, m.db_instance)
520         assert seen["id"] == "cj2"
521         done = db.get_job("cj2")
522         assert done["status"] == "done"
523         assert done["result"]["download_url"] == "/api/highlights/vidX/r.mp4"
524
525
526     def test_create_job_persists_and_reconstructs_compute(tmp_path):
527         """v8: the compute-picker choice must round-trip create_job ? get_job ?
_job_to_request so the
528         worker honours it across the async submit?reconstruct boundary. NULL ? server
default (the
529         pre-picker default-OFF behaviour). Regression guard for the dropped-across-boundary
bug."""
530         db = Database(str(tmp_path / "jobs.db"))
531         db.create_job("j-cloud", "gs://b/clip.mp4", compute="cloud")
532         db.create_job("j-default", "/v.mp4") # no compute ? NULL ? server default
533         assert db.get_job("j-cloud")["compute"] == "cloud"
534         assert db.get_job("j-default")["compute"] is None
535         assert m._job_to_request(db.get_job("j-cloud")).compute == "cloud"
536         assert m._job_to_request(db.get_job("j-default")).compute is None
537
538
539     # --- crash recovery -----

```

56. assistant (2026-07-08T05:18:30.420Z)

I have everything I need. This is a large, coherent change. Let me track it with a task list, then implement ? starting with the config layer.

57. user (2026-07-08T05:18:32.152Z)

context OK

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

58. assistant (2026-07-08T05:18:56.314Z)

Shell cwd resets each call (I'll use absolute worktree paths `C:/kg-wt521/...` everywhere). Let me read the worktree copies of files I'll edit (to register them for editing) ? starting with the config `\_\_init\_\_` which I haven't seen yet.

59. user (2026-07-08T05:18:56.384Z)

```
1      """Configuration package.
2
3      Public API:
4          from backend.config import load_config, AppConfig, Config
5
6      This package is the canonical place for all configuration access.
7      """
8
9      import os
10
11     from .loader import (
12         LIGHT_DEFAULT_CONFIG,
13         load_config,
14         save_default_config,
15         write_light_default_config,
16     )
17     from .models import ( # re-export common ones
18         AppConfig,
19         Config,
20         DeploymentConfig,
21         IndexingConfig,
22         StitchingConfig,
23         StorageConfig,
24         ValidationConfig,
25     )
26     from .presets import PROFILE_PRESETS, probe_cuda_available, suggest_profile
27     from .startup import (
28         StartupCheck,
29         StartupCheckError,
30         enforce_startup_checks,
31         run_startup_checks,
32     )
33
34
35     def default_db_path(config_path: str = "config.json") -> str:
36         """Resolve the DB path from the typed config ? the single shared helper.
37
38         Replaces the copy-pasted ``_default_db_path()`` in ``run_eval.py`` and
39         ``run_regression.py``.
40         """
41         cfg = load_config(config_path)
42         return os.path.join(cfg.output.output_dir, cfg.output.db_name)
43
44
45     __all__ = [
46         "load_config",
47         "save_default_config",
48         "write_light_default_config",
49         "LIGHT_DEFAULT_CONFIG",
50         "default_db_path",
51         "AppConfig",
52         "Config",
```

```

53         "DeploymentConfig",
54         "IndexingConfig",
55         "StitchingConfig",
56         "StorageConfig",
57         "ValidationConfig",
58         "PROFILE_PRESETS",
59         "suggest_profile",
60         "probe_cuda_available",
61         "StartupCheck",
62         "StartupCheckError",
63         "enforce_startup_checks",
64         "run_startup_checks",
65     ]
66

```

60. assistant (2026-07-08T05:18:59.523Z)

Now let me implement. First, the new placement resolver module:

61. assistant (2026-07-08T05:19:13.203Z)

```

"""Phase ? machine placement resolution (#521).

```

A *phase placement* answers one question per pipeline phase ? *where does the heavy compute for this phase run?* ? with a config knob (`deployment.phase_placement`) instead of code. The three placeable phases and their values:

```

validation ? control_plane | cloud_run_l4
segment    ? control_plane | cloud_run_l4 (a.k.a. "detect")
compile    ? control_plane | cloud_run_l4 (a.k.a. "stitch")

```

`control_plane` runs the phase IN-PROCESS on the (CPU-only) control-plane service;

`cloud_run_l4`

dispatches it to the GPU/NVENC Cloud Run worker Job (the same L4 detection already uses). This module

is the SINGLE source of truth so `backend.orchestration` (segment + compile) and the worker never

re-derive the mapping.

Default = today's behaviour (pure / additive). When a phase's knob is unset (`None`) the resolver reproduces the pre-#521 placement:

```

validation ? control_plane (the authoritative gate,
#512/#513)
segment    ? cloud_run_l4 IFF deployment.compute_target == "cloud_run", else control_plane
compile    ? control_plane (the in-process ffmpeg
stitch)

```

so a config with no `phase_placement` block behaves byte-identically to before this module existed.

```

"""

```

```

from __future__ import annotations

```

```

from typing import Any

```

Phase identities (kept as constants so callers don't stringly-type the phase names).

```

PHASE_VALIDATION = "validation"

```

```

PHASE_SEGMENT = "segment"

```

```

PHASE_COMPILE = "compile"

```

```

PHASES: tuple[str, ...] = (PHASE_VALIDATION, PHASE_SEGMENT, PHASE_COMPILE)

```

Placement targets.

```
PLACE_CONTROL_PLANE = "control_plane"
PLACE_CLOUD_L4 = "cloud_run_l4"
PLACEMENTS: tuple[str, ...] = (PLACE_CONTROL_PLANE, PLACE_CLOUD_L4)
```

```
def resolve_phase_placement(cfg: Any, phase: str) -> str:
    """Where does ``phase`` run ? ``"control_plane"`` or ``"cloud_run_l4"``?
```

Precedence: an explicit ``deployment.phase\_placement.<phase>`` wins; otherwise the legacy default

(mirrors the pre-#521 behaviour, see the module docstring). Never raises ? an absent
``deployment`` / ``phase\_placement`` degrades to the legacy default.
"""

```
deployment = getattr(cfg, "deployment", None)
placement = getattr(deployment, "phase_placement", None)
explicit = getattr(place, phase, None) if placement is not None else None
if explicit:
    return str(explicit)
# Legacy defaults ? reproduce today's placement exactly.
if phase == PHASE_SEGMENT:
    compute_target = (getattr(deployment, "compute_target", "") or "").strip()
    return PLACE_CLOUD_L4 if compute_target == "cloud_run" else PLACE_CONTROL_PLANE
# validation + compile default to the control-plane (their pre-#521 home).
return PLACE_CONTROL_PLANE
```

62. user (2026-07-08T05:19:13.235Z)

File created successfully at: C:\kg-wt521\backend\config\placement.py (file state is current in your context ? no need to Read it back)

63. assistant (2026-07-08T05:19:18.274Z)

Now edit the config models. Let me register the worktree file with targeted reads, then edit.

64. user (2026-07-08T05:19:19.264Z)

```
354     class StitchingConfig(BaseModel):
355         begin_padding: float = 0.5
356         end_padding: float = 0.5
357         use_cache: bool = False
358         min_shot_count: int = 1
359         # Long-rally reel filter (seconds; 0 = OFF). Drop detected rally pairs whose
duration
360         # (end ? start) is below this from the highlight reel, BEFORE compiling. A cleaner
361         # reel-selection knob than `min_shot_count`: shot_count is unlabeled / "Unknown" on
the
362         # fully-local no-AI path, whereas duration is derived from the rally's own
start/end. The
363         # local detector over-fires and its false positives are mostly SHORT (motion blips,
364         # background-court fragments), so this keeps the reel to the eye-catching long
rallies.
365         # OFF by default ? the detector output is unchanged, only which rallies reach the
reel.
366         min_rally_duration: float = Field(default=0.0, ge=0.0)
367         watermark: WatermarkConfig = Field(default_factory=WatermarkConfig)
368
```

65. user (2026-07-08T05:19:19.743Z)

```
454     class DeploymentConfig(BaseModel):
455         """Deployment profile ? one switch that selects a coherent bundle of compute +
storage
456         knobs (COMPUTE_DECOUPLED_SERVING M1, ``docs/COMPUTE_DECOUPLED_SERVING/README.md``
```

§4.3).

```

457
458     **Pure / additive:** the empty default profile (``"``) applies NO preset, so the
pipeline
459     behaves exactly as before this section existed. A profile is opt-in (``config.json``
460     ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); auto-detect only ever
461     *suggests* one (D15 ? cloud spend is never entered silently). Nothing in the core
reads
462     these fields yet (M1): selecting a profile works purely by pre-setting the EXISTING
knobs
463     (``indexing.detector_impl`` / ``skip_ai_handoff``, ``storage.backend``) in the
loader.
464
465     Valid values mirror ``backend/config/presets.py`` (parity-pinned by a unit test)."""
466
467     # "" = no profile (today's manual knobs). Others apply presets.PROFILE_PRESETS.
468     profile: Literal["", "truly-local", "cloud-serving", "cpu-mock"] = ""
469     # Where the core runs. Set by the profile preset; overridable in config.json. Inert
in M1
470     # (recorded for observability; consumed by the ComputeBackend seam from M2+).
471     compute_target: Literal["", "local_gpu", "cloud_run", "cpu_mock"] = ""
472     # Cloud-serving throughput cap (#466): how many DETECT jobs the control-plane drain
may dispatch +
473     # bucket-poll CONCURRENTLY when compute_target=="cloud_run". Each is an I/O-bound
Cloud Run Job
474     # execution (the GPU work is on ephemeral L4s, WASB-only has no Gemini), so ? unlike
a truly-local
475     # single box ? they need not serialize. Bounds concurrent PAID L4 executions. Non-
cloud_run targets
476     # ignore this and stay serial (max_workers=1), preserving the truly-local GPU/Gemini
serialization.
477     max_concurrent_remote_jobs: int = Field(default=4, ge=1, le=64)
478     # Poll budget for one dispatched remote job (CONTROL_PLANE_DURABLE_REBUILD §4.3).
Must sit
479     # ABOVE the worker Job's own --task-timeout (3600s) so the worker's timeout stays
the
480     # authority and the dispatcher's one-shot post-timeout marker rescue still applies ?
the old
481     # 1800s ExecutionPolicy default undercut real 1080p60 jobs (~22 min detect +
queue/cold-start,
482     # observed on the 2026-07-05 owner CUJ) and could abandon a slow-but-healthy paid
GPU run.
483     remote_poll_max_wait_sec: float = Field(default=3900.0, gt=0)

```

66. assistant (2026-07-08T05:19:23.358Z)

Now the StitchingConfig encoder field:

67. user (2026-07-08T05:19:29.794Z)

The file C:\kg-wt521\backend\config\models.py has been updated successfully. (file state is current in your context ? no need to Read it back)

68. assistant (2026-07-08T05:19:33.470Z)

Now add the `PhasePlacementConfig` model and wire it into `DeploymentConfig` with the encoder knob and the footgun guard:

69. user (2026-07-08T05:19:40.671Z)

The file C:\kg-wt521\backend\config\models.py has been updated successfully. (file state is current in your context ? no need to Read it back)

70. user (2026-07-08T05:19:53.301Z)

The file C:\kg-wt521\backend\config\models.py has been updated successfully. (file state is current in your context ? no need to Read it back)

71. assistant (2026-07-08T05:19:56.153Z)

Now export the resolver and new model from the config package:

72. user (2026-07-08T05:19:59.936Z)

The file C:\kg-wt521\backend\config__init__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

73. user (2026-07-08T05:20:04.846Z)

The file C:\kg-wt521\backend\config__init__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

74. assistant (2026-07-08T05:20:08.941Z)

Now set compile?cloud_run_l4 in the cloud-serving preset:

75. user (2026-07-08T05:20:09.549Z)

```
51         "cloud-serving": {
52             "deployment": {"profile": "cloud-serving", "compute_target": "cloud_run"},
53             # Linux L4 (cuda); Gemini handoff ON until the owned model clears the #172 floor
54             # wasb GPU-feed tuning: the parity defaults (batch 8, no prefetch) left the paid
55             # <1% utilized on the first real serving run (2026-07-03) ? the pipeline was
56             # CPU-feed-bound. Serving opts into the measured fast values; local/WSL
57             "indexing": {
58                 "detector_impl": "native",
59                 "skip_ai_handoff": False,
60                 # GPU-resident NVDEC decode (#506/#507): validated on a real L4 (golden gate
61                 # F1 0.574 ? baseline 0.582 @ SERVED) AND on the live Cloud Run L4 (the
62                 # trajectory reproduced the VM run ? 26479 ? 26481 visible on GX010128),
63                 # out + canaried (DEPLOY_REPORT_cloud-serving-gpu-resident-
64                 # nvdec_2026-07-07). This makes
65                 # nvdec the code-level serving default so it survives an image rebuild; the
66                 # WASB_DECODE_BACKEND env was the reversible rollout lever. cuda-only
67                 "wasb": {
68                     "batch_size": 64,
69                     "prefetch_batches": 4,
70                     "decode_backend": "nvdec_resident",
71                 },
72                 "storage": {"backend": "gcs"},
73             },
```

76. user (2026-07-08T05:20:20.983Z)

The file C:\kg-wt521\backend\config\presets.py has been updated successfully. (file state is current in your context ? no need to Read it back)

77. assistant (2026-07-08T05:20:30.706Z)

Now the compute layer. Let me add `CompileJobSpec` and `run\_compile` to the base:

78. user (2026-07-08T05:20:31.395Z)

```
52
53
54     @dataclass(frozen=True)
55     class ComputeResult:
56         """The terminal state of a dispatched job. ``returncode`` mirrors the worker's exit
code
57         (0 ok / 2 validation-or-config / 3 segmenter-failure), plus **4 = remote dispatch
timeout or an
58         empty/unreadable completion marker** (the cloud path only). ``video_id`` is known
once the job's
59         done-marker is read; ``outputs_uri`` is where ``intervals.csv`` / ``report.json``
landed."""
60
61         returncode: int
62         outputs_uri: str = ""
63         video_id: str = ""
64         report: dict = field(default_factory=dict)
65         # The backend's execution identity for PROVENANCE (e.g. the Cloud Run execution
name) ? surfaced
66         # on the job result so "the cloud path actually ran" is independently verifiable
(gcloud run jobs
67         # executions describe <execution>). Empty for the in-process path / backends without
one.
68         execution: str = ""
69
70
71     class ComputeBackend(ABC):
72         """Backend-agnostic compute. ``run_infer`` dispatches ONE inference job and returns
only once
73         its outputs are durably available (bucket-poll for cloud), so a caller treats local
and cloud
74         execution identically."""
75
76         @abstractmethod
77         def run_infer(
78             self,
79             spec: ComputeJobSpec,
80             *,
81             on_wait: "Callable[[float], None] | None" = None,
82         ) -> ComputeResult:
83             """Run ``spec`` to completion and return its ``ComputeResult``.
84
85             ``on_wait(elapsed_sec)`` (optional) is invoked periodically while the backend is
86             still WAITING on the remote work ? the caller's live-progress heartbeat (#482),
so a
87             long GPU run doesn't look frozen. Implementations without a wait loop may ignore
it."""
88
```

79. user (2026-07-08T05:20:40.427Z)

The file C:\kg-wt521\backend\compute\base.py has been updated successfully. (file state is current in your context ? no need to Read it back)

80. user (2026-07-08T05:20:48.301Z)

The file C:\kg-wt521\backend\compute\base.py has been updated successfully. (file state is current in your context ? no need to Read it back)

81. assistant (2026-07-08T05:20:50.883Z)

Now export `CompileJobSpec` from the compute package:

82. user (2026-07-08T05:20:54.593Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

83. user (2026-07-08T05:20:57.796Z)

```
16     from backend.compute.base import (
17         ComputeBackend,
18         ComputeJobSpec,
19         ComputeResult,
20         RemoteExecutionFailed,
21     )
22     from backend.infrastructure.execution_policy import DEFAULT_POLICY, ExecutionPolicy
23
24     __all__ = [
25         "ComputeBackend",
26         "ComputeJobSpec",
27         "ComputeResult",
28         "RemoteExecutionFailed",
29         "get_compute_backend",
30     ]
```

84. user (2026-07-08T05:21:04.505Z)

The file C:\kg-wt521\backend\compute__init__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

85. assistant (2026-07-08T05:21:09.444Z)

Now the `CloudRunCompute` ? extract a shared marker-await helper and add `run\_compile`. Let me register and read the worktree file's relevant regions:

86. user (2026-07-08T05:21:10.281Z)

```
244     def run_infer(
245         self,
246         spec: ComputeJobSpec,
247         *,
248         on_wait: "Callable[[float], None] | None" = None,
249     ) -> ComputeResult:
250         out_root = spec.out_uri.rstrip("/")
251         token = self._token or uuid.uuid4().hex
252         done_uri = f"{out_root}/_dispatch/{token}.done"
253         argv: List[str] = [
254             "infer",
255             "--video-uri",
256             spec.video_uri,
257             "--out-uri",
258             spec.out_uri,
259             "--done-uri",
260             done_uri,
261         ]
262         if spec.segmenter:
263             argv += ["--segmenter", spec.segmenter]
264         if spec.skip_validation:
265             # The control-plane already validated + only dispatches on PASS ? tell the
worker not to
266             # re-decode and re-validate (redundant, and its config may resolve different
thresholds).
267             argv.append("--skip-validation")
268         for kv in (*spec.config_overrides, *self._container_overrides):
269             argv += ["--set", kv]
270
271         execution = self._client.run_job(
```

```

272         self._job_name, argv, region=self._region, project=self._project
273     )
274     LOG.info(
275         "cloud-run: dispatched job=%s exec=%s; bucket-polling %s",
276         self._job_name,
277         execution,
278         done_uri,
279     )
280
281     try:
282         marker = poll_until(
283             lambda: self._poll_check(done_uri, str(execution or "")),
284             self._policy,
285             label="cloud-run-infer",
286             logger=LOG,
287             # Live heartbeat (#482): each still-waiting tick reaches the caller so
288             # job.progress moves during the whole GPU run instead of freezing.
289             on_tick=on_wait,
290         )
291     except PolicyTimeout as timeout_exc:
292         # RemoteExecutionFailed (the fast negative path in _poll_check) is NOT a
PolicyTimeout,
293         # so it propagates past here uncaught ? the dispatch handler maps it to a
distinct,
294         # clearly-labelled failure instead of the ~65-min opaque timeout this branch
handles.
295         marker = self._handle_poll_timeout(done_uri, str(execution), timeout_exc)
296         video_id = str(marker.get("video_id", ""))
297         # A MISSING returncode defaults to FAILURE (1), never success ? a present-but-
unreadable /
298         # garbled marker must not masquerade as rc=0. (The worker always writes both
fields.)
299         rc = int(marker.get("returncode", 1))
300         if not video_id:
301             LOG.warning(
302                 "cloud-run: job=%s completion marker present but empty/unreadable (no "
303                 "video_id) ? treating as failure",
304                 self._job_name,
305             )
306             rc = rc or 1
307         outputs_uri = f"{out_root}/outputs/{video_id}/" if video_id else ""
308         report = self._read_json(f"{outputs_uri}report.json") if outputs_uri else {}
309         LOG.info(
310             "cloud-run: job=%s done (video_id=%s rc=%d)", self._job_name, video_id, rc
311         )
312         return ComputeResult(
313             returncode=rc,
314             outputs_uri=outputs_uri,
315             video_id=video_id,
316             report=report,
317             execution=str(execution or ""),
318         )

```

87. user (2026-07-08T05:21:18.458Z)

The file C:\kg-wt521\backend\compute\cloud_run.py has been updated successfully. (file state is current in your context ? no need to Read it back)

88. assistant (2026-07-08T05:21:20.895Z)

Now add `\_await\_marker` and `\_run\_compile` after `\_run\_infer` (before `\_poll\_check`):

89. user (2026-07-08T05:21:36.491Z)

The file C:\kg-wt521\backend\compute\cloud_run.py has been updated successfully. (file state is current in your context ? no need to Read it back)

90. assistant (2026-07-08T05:21:39.404Z)

Now import `CompileJobSpec` into cloud_run.py:

91. user (2026-07-08T05:21:41.185Z)

The file C:\kg-wt521\backend\compute\cloud_run.py has been updated successfully. (file state is current in your context ? no need to Read it back)

92. assistant (2026-07-08T05:21:47.282Z)

Now the VideoCompiler NVENC encoder support. Let me register the worktree file and edit the constructor, add encoder selection, and use it in the trim:

93. user (2026-07-08T05:21:48.213Z)

```
27     class VideoCompiler:
28         def __init__(
29             self,
30             output_dir: str,
31             end_padding: float = 1.0,
32             begin_padding: float = 0.5,
33             watermark: Optional[Any] = None,
34         ):
35             self.output_dir = output_dir
36             self.end_padding = end_padding
37             self.begin_padding = begin_padding
38             # Optional branding overlay. Accepts a WatermarkConfig model or a plain dict;
39             # normalized to a dict (or None) so this module stays config-package-agnostic.
40             self.watermark = self._normalize_watermark(watermark)
41             os.makedirs(output_dir, exist_ok=True)
42
```

94. user (2026-07-08T05:21:56.522Z)

The file C:\kg-wt521\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

95. assistant (2026-07-08T05:22:03.969Z)

Now add the encoder-selection logic. First a module-level probe cache, then the selection method:

96. user (2026-07-08T05:22:13.977Z)

The file C:\kg-wt521\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

97. assistant (2026-07-08T05:22:16.596Z)

Now add the `\_select\_encoder` method (before `\_trim\_one\_segment`) and thread the resolved codec args through the trim:

98. user (2026-07-08T05:22:27.278Z)

The file C:\kg-wt521\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

99. assistant (2026-07-08T05:22:29.726Z)

Now use the codec args in the `base\_cmd`:

100. user (2026-07-08T05:22:34.809Z)

The file C:\kg-wt521\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

101. assistant (2026-07-08T05:22:36.849Z)

Now resolve the encoder once in `compile\_highlights` and pass it to each trim:

102. user (2026-07-08T05:22:43.519Z)

The file C:\kg-wt521\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

103. assistant (2026-07-08T05:22:49.769Z)

Now the core: `backend/orchestration.py`. Let me register the worktree file and add the placement import:

104. user (2026-07-08T05:22:50.322Z)

```
21     import csv
22     import json
23     import logging
24     import math
25     import os
26     import re
27     import tempfile
28     import threading
29     from typing import TYPE_CHECKING, Any, Dict, List, Optional, Tuple
30     from urllib.parse import quote
31
32     from backend import storage
33     from backend.config import AppConfig, StitchingConfig
34     from backend.pipeline.segmenters.base import segmenter_registry
35     from backend.pipeline.stitcher.watermark_i18n import localize_watermark
36     from backend.pipeline.validators.base import ValidationResult, registry
37     from backend.results import Failure
38     from backend.utils.paths import output_base, safe_output_filename
```

105. user (2026-07-08T05:22:54.362Z)

The file C:\kg-wt521\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

106. assistant (2026-07-08T05:22:58.513Z)

Now add the `\_will\_detect\_on\_cloud` predicate after `\_cloud\_available`:

107. user (2026-07-08T05:23:05.745Z)

The file C:\kg-wt521\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

108. assistant (2026-07-08T05:23:10.786Z)

Now update `\_maybe\_dispatch\_remote`'s default target branch to consult segment placement:

109. user (2026-07-08T05:23:16.434Z)

The file C:\kg-wt521\backend\orchestration.py has been updated successfully. (file state is

current in your context ? no need to Read it back)

110. assistant (2026-07-08T05:23:19.566Z)

Now the `skip\_validation` derivation in the `ComputeJobSpec`:

111. user (2026-07-08T05:23:24.732Z)

The file C:\kg-wt521\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

112. assistant (2026-07-08T05:23:28.496Z)

Now wire the validation-delegation into `\_execute\_processing`. Let me read the current validation block region:

113. user (2026-07-08T05:23:29.215Z)

```
795         storage.cleanup_acquired_local_input(
796             req.video_path, local_video, getattr(config, "storage", None)
797         )
798         raise
799     try:
800         existing_video = db.get_video(video_id)
801     except Exception:
802         storage.cleanup_acquired_local_input(
803             req.video_path, local_video, getattr(config, "storage", None)
804         )
805         raise
806     was_reingest = existing_video is not None
807
808     # typed AppConfig: attribute access for the default segmenter (with safe fallback)
809     default_seg = getattr(
810         getattr(config, "indexing", None), "default_segmenter", "motion"
811     )
812     seg_name = req.segmenter or default_seg
813     segmenter_actual = None
814     segmentation_ran = False
815     val_results: List[Any] = []
816     val_context: Dict[str, Any] = {}
817     # Compute-path PROVENANCE (the validation method): which backend actually ran the
DETECT ?
818     # "in_process" once the local segmenter is reached, set on the response as
"cloud_run" by
819     # _maybe_dispatch_remote when it dispatches. None ? nothing ran (validation/config
fail).
820     compute_actual: Optional[str] = None
821     response: Optional[Dict[str, Any]] = None
822     try:
823         if (
824             existing_video
825             and existing_video.get("status") == "processed"
826             and not req.force
827         ):
828             cached_segments = db.query_play_segments(video_id, {})
829             if cached_segments:
830                 logger.info(
831                     "[API] Reusing cached segmentation for video_id=%s; force=true
bypasses cache.",
832                     video_id,
833                 )
834             response = _cached_process_payload(
835                 video_id,
836                 existing_video.get("validation_results", []),
```

```

837             cached_segments,
838             getattr(req, "compute", None),
839         )
840         return response
841
842         # 1. Run pluggable validators (with-context variant surfaces ffprobe_meta for
the report) ?
843         #     UNLESS this exact content already has a cached verdict under the current
validator config.
844         notify("validating")
845         # VALIDATION-RESULT CACHE: video_id is the MD5 of the file bytes, so an
identical-content
846         # reupload keys the same row. When that content was ALREADY validated under the
SAME validator
847         # configuration, reuse the stored verdict instead of re-decoding the clip
through the whole
848         # validator suite ? minutes on a heavy 4K clip on the CPU-only control-plane,
which the
849         # 2026-07-07 CUJ re-paid on every retry after a segmentation failure (validation
passed but
850         # detection failed ? no cached segments, so the whole-pipeline cache above
misses and we land
851         # here). The fingerprint folds in the validation.* config + sport + validator
set + a schema
852         # version, so ANY config change misses and re-validates (never serve a stale
PASS/FAIL).
853         # ``force=true`` bypasses the reuse and re-writes a fresh verdict below (mirrors
how it bypasses
854         # the whole-pipeline cache). Composes with #513 (the worker skips its own re-
validation on a
855         # cloud dispatch), so a post-failure reupload reaches segmentation in ~0
validation time.
856         val_fingerprint = registry.config_fingerprint(config)
857         # The whole read ? fingerprint-match ? replay is guarded: ANY cache fault
(unreadable row,
858         # corrupt or schema-incompatible stored results) degrades to a real validation
rather than
859         # failing the job. The cache is a pure speedup ? never a new failure mode.
860         reused_validation = False
861         if not req.force:
862             try:
863                 cached = db.get_validation_cache(video_id)
864                 if cached is not None and cached.get("fingerprint") == val_fingerprint:
865                     # Replay the stored ValidationResult list; val_context stays {} (no
fresh
866                     # ffprobe_meta ? the report's media block degrades gracefully to
size-only, the
867                     # deliberate cost of skipping the decode). Segmentation reads
neither of these.
868                     val_results = [
869                         ValidationResult(**r) for r in cached.get("results", [])
870                     ]
871                     reused_validation = True
872                     logger.info(
873                         "[API] Reusing cached validation verdict (passed=%s) for
video_id=%s; "
874                         "validator config unchanged ? skipping the validator suite. "
875                         "force=true re-validates.",
876                         cached.get("passed"),
877                         video_id,
878                     )
879             except Exception: # noqa: BLE001 ? any cache read/replay fault degrades to
a real validation
880                 reused_validation = False
881                 logger.warning(

```

```

882             "[API] validation-cache reuse failed for video_id=%s ? re-
validating.",
883             video_id,
884             exc_info=True,
885         )
886     if not reused_validation:
887         logger.info(f"Running validations for {req.video_path}...")
888         val_results, val_context = registry.run_all_with_context(local_video,
config)
889         failures = [r for r in val_results if not r.passed]
890     if not reused_validation:
891         # Persist the verdict keyed by (content, config) so the NEXT identical
reupload skips the
892         # decode. Best-effort (the writer swallows + logs its own faults); a miss
just re-validates.
893         db.set_validation_cache(
894             video_id,

```

114. user (2026-07-08T05:23:33.726Z)

```

886     if not reused_validation:
887         logger.info(f"Running validations for {req.video_path}...")
888         val_results, val_context = registry.run_all_with_context(local_video,
config)
889         failures = [r for r in val_results if not r.passed]
890     if not reused_validation:
891         # Persist the verdict keyed by (content, config) so the NEXT identical
reupload skips the
892         # decode. Best-effort (the writer swallows + logs its own faults); a miss
just re-validates.
893         db.set_validation_cache(
894             video_id,
895             val_fingerprint,
896             not failures,
897             [r.model_dump() for r in val_results],
898         )
899
900     if failures:
901         # add_video is an UPSERT (no cascade) ? a failed re-validation only updates
902         # status; it no longer destroys the video's prior good segments.
903         db.add_video(
904             video_id,
905             req.video_path,
906             "failed",
907             [r.model_dump() for r in val_results],
908         )
909         response = {
910             "video_id": video_id,
911             "status": "failed",
912             "message": _validation_failure_message(failures),
913             "results": [r.model_dump() for r in val_results],
914         }
915         return response
916
917     # 2. Run segmenter & indexer
918     logger.info("[API] Video passed checks. Segmenting and indexing...")
919     db.add_video(
920         video_id, req.video_path, "processed", [r.model_dump() for r in val_results]
921     )
922
923     segmenter_cls = segmenter_registry.get(seg_name)
924     if not segmenter_cls:
925         # Submit-time validation already 400s unknown names; this is the defensive

```

115. assistant (2026-07-08T05:24:00.876Z)

I'll gate the validator suite on a `validate\_on\_control\_plane` flag (delegating to the worker when validation is placed on the L4). Let me make targeted edits:

116. user (2026-07-08T05:24:10.406Z)

The file C:\kg-wt521\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

117. user (2026-07-08T05:24:16.669Z)

The file C:\kg-wt521\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

118. assistant (2026-07-08T05:24:21.219Z)

Now the compile dispatch. Let me read the current `\_execute\_compile` region in the worktree (shifted by earlier edits):

119. user (2026-07-08T05:24:21.833Z)

```
1155         return f"{root.rstrip('/')}/highlights/{video_id}/{name}"
1156
1157
1158     def _storage_settings(cfg) -> Optional[Dict[str, Any]]:
1159         """The active storage settings dict (per-backend config for the storage layer ? S3
endpoint/region,
1160         GCS project/client/signed_url_ttl), or None. Threaded into storage calls so a non-
default backend
1161         isn't silently constructed as a fresh default."""
1162         sc = getattr(cfg, "storage", None)
1163         return sc.model_dump() if sc is not None else None
1164
1165
1166     def signed_reel_url(cfg, video_id: str, name: str) -> Optional[str]:
1167         """A short-lived signed URL for a reel mirrored to the output bucket, so
/api/highlights can
1168         307-redirect off the app's response path (bypassing Cloud Run's ~32 MiB cap + the
ephemeral local
1169         /tmp; #463). None ? no bucket copy ? serve the local file. Uses the active storage
config for BOTH
1170         the existence check and the signing (extracted here to keep backend/main.py lean ?
P23/P24)."""
1171         reel_uri = reel_object_uri(cfg, video_id, name)
1172         if not reel_uri:
1173             return None
1174         settings = _storage_settings(cfg)
1175         if storage.exists(reel_uri, config=settings):
1176             return storage.signed_url(reel_uri, config=settings, method="GET")
1177         return None
1178
1179
1180     def _execute_compile(cfg, db, job: Dict[str, Any], progress=None) -> Dict[str, Any]:
1181         """Worker handler for a ``kind='compile'`` job (#329): re-resolve the segments,
stitch the reel,
1182         and return the {status, filepath, download_url} result the SPA polls for. Returns a
1183         ``status='failed'`` result for an EXPECTED problem (missing video / no segments / a
stitch error)
1184         so the SPA always gets a clean verdict instead of a worker-crash error."""
1185         import backend.main as _main # call-time seam: VideoCompiler stays patchable on
backend.main
1186
1187         notify = progress or (lambda stage: None)
1188         params = json.loads(job.get("params") or "{}")
1189         video_id = job.get("video_id") or ""
```



```

1190     segment_ids = params.get("segment_ids") or []
1191     locale = params.get("locale")
1192
1193     notify("resolving")
1194     try:
1195         video, kept, out_name = _resolve_compile(
1196             db, cfg, video_id, segment_ids, params.get("output_filename", "")
1197         )
1198     except _CompileError as e:
1199         return {"status": "failed", "message": e.detail, "video_id": video_id}
1200
1201     stitching = getattr(cfg, "stitching", None) or StitchingConfig()
1202     output_dir = (
1203         getattr(getattr(cfg, "output", None), "output_dir", "output") or "output"
1204     )
1205     # Localize the branding watermark to the UI locale the SPA rendered in (recorded on
1206     # the job);
1207     # unknown/absent locale keeps the configured default unchanged.
1208     localized_watermark = localize_watermark(stitching.watermark, locale)
1209     compiler = _main.VideoCompiler(
1210         output_dir,
1211         begin_padding=stitching.begin_padding,
1212         end_padding=stitching.end_padding,
1213         watermark=localized_watermark,
1214     )
1215     time_pairs = [(s["start_time"], s["end_time"]) for s in kept]
1216
1217     notify("stitching")
1218     local_src = ""
1219     try:
1220         # The stored `filepath` is the SOURCE ref (gs://, s3://, gdrive://, or a local
1221         # path). Resolve to
1222         # a LOCAL file the stitcher can read ? identity for a local path, fetch-to-
1223         # scratch for a URI (the
1224         # same seam the process path uses); without it, compiling a bucket-ingested
1225         # video can't open gs://.
1226         local_src = storage.acquire_local_input(
1227             video["filepath"], getattr(cfg, "storage", None)
1228         )
1229         # The ffmpeg stitch is CPU-heavy ? serialize it so a raised drain worker count
1230         # (for remote
1231         # dispatch) doesn't run multiple compiles at once on the box (see
1232         # _LOCAL_COMPUTE_LANE).
1233         with _LOCAL_COMPUTE_LANE:
1234             out_path = compiler.compile_highlights(local_src, time_pairs, out_name)
1235     except Exception as e: # noqa: BLE001 ? a stitch/ffmpeg failure is a clean job
1236         # FAILURE, not a crash
1237         logger.error(
1238             "[compile] stitching failed for video_id=%s: %s", video_id, e, exc_info=True
1239         )
1240     return {
1241         "status": "failed",
1242         "message": f"Highlight stitching failed: {e}",
1243         "video_id": video_id,
1244     }
1245     finally:
1246         if local_src:
1247             storage.cleanup_acquired_local_input(
1248                 video["filepath"], local_src, getattr(cfg, "storage", None)
1249             )
1250
1251     # #463: mirror the reel to the output bucket so /api/highlights can 307 to a signed
1252     # URL ? Cloud Run
1253     # caps a buffered response at ~32 MiB (a real reel is bigger) and the local /tmp
1254     # copy dies when the

```

```

1246         # instance recycles. A mirror-upload failure must NOT fail an otherwise-good stitch
(the endpoint
1247         # still falls back to the local file on the same instance).
1248         reel_uri = reel_object_uri(cfg, video_id, out_name)
1249         if reel_uri:
1250             try:
1251                 notify("uploading reel")
1252                 # Thread the active storage settings so the configured backend (S3
endpoint/region, GCS
1253                 # project/client) is used, not a fresh default.
1254                 storage.put(out_path, reel_uri, config=_storage_settings(cfg))
1255             except Exception as e: # noqa: BLE001 ? best-effort mirror; local stitch
already succeeded
1256                 logger.warning("[compile] could not mirror reel to %s: %s", reel_uri, e)
1257
1258         # `filepath` is server-local (not browser-reachable); also return the download URL
so the SPA needn't
1259         # hand-build it (P6 / PER_RALLY_SELECTION.md §4.3).
1260         download_url = f"/api/highlights/{quote(video_id)}/{quote(out_name)}"
1261         notify("finished")
1262         return {
1263             "status": "success",
1264             "filepath": out_path,

```

120. user (2026-07-08T05:24:24.779Z)

```

1262         return {
1263             "status": "success",
1264             "filepath": out_path,
1265             "download_url": download_url,
1266             "video_id": video_id,
1267         }
1268

```

121. assistant (2026-07-08T05:25:23.516Z)

Now I'll refactor `\_execute\_compile` ? extract the shared `\_stitch\_reel` core, add the placement branch, and the `\_dispatch\_compile\_remote` handler. Let me replace the whole function:

122. user (2026-07-08T05:26:25.553Z)

The file C:\kg-wt521\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

123. assistant (2026-07-08T05:26:38.297Z)

Let me check the `backend.main` re-export block to add the new orchestration symbols for addressability:

124. user (2026-07-08T05:26:39.177Z)

```

226     )
227
228     # The processing/compile ORCHESTRATION (hash?validate?segment?report pipeline, Cloud Run
229     # dispatch+ingest, compile resolution + worker handler) now lives in
backend.orchestration
230     # (P23 god-module split; re-cut of #390). Re-exported here so the names stay addressable
on
231     # backend.main ? `import backend.main as m; m._execute_processing(...)` , the job
worker's
232     # call-time `_main._execute_compile(...)` seam, and monkeypatch.setattr(backend.main,
...)
233     # all keep working unchanged.
234     from backend.orchestration import ( # noqa: F401 (re-exports: names stay addressable

```

```

on backend.main)
235     _cloud_available,
236     _CompileError,
237     _execute_compile,
238     _execute_processing,
239     _finalize_reingest,
240     _ingest_remote_segments,
241     _maybe_dispatch_remote,
242     _remote_input_size_bytes,
243     _report_config_for_input,
244     _resolve_compile,
245     read_path_rehydrate,
246     signed_reel_url,
247 )
248
249
250 # Legacy thin wrapper for any remaining direct calls during transition.
251 # New code should import load_app_config / AppConfig from backend.config directly.
252 def load_config(config_path: str = "config.json") -> Dict[str, Any]:
253     cfg = load_app_config(config_path)
254     return cfg.model_dump(mode="json")
255
256
257 # --- API Models ---

```

125. user (2026-07-08T05:26:45.055Z)

The file C:\kg-wt521\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

126. assistant (2026-07-08T05:27:35.071Z)

Now the worker `compile` subcommand. Let me register the worktree worker file and add `cmd\_compile` + helpers after `cmd\_infer`:

127. user (2026-07-08T05:27:36.037Z)

```

415     if getattr(args, "done_uri", None):
416         _write_done_marker(
417             args.done_uri, video_id, 0, len(segments), status, storage_cfg, scratch
418         )
419     return 0
420
421
422 #: Video extensions accepted under <corpus>/videos/ (mirrors the labels/ '.csv' filter
so a
423 #: per-video sidecar ? .srt/.json/thumbnail ? can't shadow the real video on a stem
collision).
424 _VIDEO_EXTS = ("mp4", ".mov", ".mkv", ".avi", ".webm", ".m4v")
425
426 #: Leading ISO-date prefix on a canonical label name (`labels/<YYYY-MM-
DD>_<name>.rallies.csv`,

```

128. user (2026-07-08T05:27:59.483Z)

The file C:\kg-wt521\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

129. assistant (2026-07-08T05:28:03.203Z)

Now add the `compile` subparser. Let me read the parser region and add it after the infer subparser:

130. user (2026-07-08T05:28:04.261Z)

```

852         metavar="k=v",
853         help="config override, e.g. --set indexing.skip_ai_handoff=true",
854     )
855     pi.add_argument(
856         "--skip-validation",
857         action="store_true",
858         help="skip the worker's re-validation: the dispatcher (control-plane) already
ran the "
859         "validators with its resolved config and only dispatches on PASS, so re-
validating here is "
860         "redundant, decodes the video again, and can contradict the gate. Default-OFF:
the direct "
861         "`worker infer` CLI still validates.",
862     )
863     pi.set_defaults(func=cmd_infer)
864
865     pt = sub.add_parser(
866         "train",
867         help="Gen-0 train-from-bucket (labels [+videos] -> trajectories -> harness)",
868     )
869     pt.add_argument(

```

131. user (2026-07-08T05:28:13.542Z)

The file C:\kg-wt521\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

132. assistant (2026-07-08T05:28:17.273Z)

Core implementation done. Let me run a quick import + lint smoke, then the existing target tests to catch regressions:

133. user (2026-07-08T05:28:22.701Z)

```

placement defaults: control_plane control_plane control_plane
imports OK
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

134. user (2026-07-08T05:28:30.795Z)

```

cloud compile: cloud_run_l4
cloud segment: cloud_run_l4
cloud validation: control_plane
GUARD OK: 1 validation error for AppConfig
deployment
  Value error, d
validation+cloud-segment OK
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

135. assistant (2026-07-08T05:28:33.876Z)

The resolver and guard work. Now let me run the existing target tests to check for regressions:

136. user (2026-07-08T05:29:43.698Z)

```

E
E     Omitting 4 identical items, use -vv to show
E     Left contains 2 more items:
E     {'cloud_stitch_encoder': 'h264_nvenc',
E       'phase_placement': {'compile': None, 'segment': None, 'validation': None}}
E     Use -v to get more diff

```

```

tests\test_config_deployment.py:395: AssertionError
===== warnings summary =====

```

```

tests/test_api_cloud_dispatch.py::test_full_api_process_to_query_via_cloud
tests/test_api_cloud_dispatch.py::test_full_api_process_to_query_via_cloud
tests/test_api_cloud_dispatch.py::test_local_override_survives_api_boundary
tests/test_api_cloud_dispatch.py::test_local_override_survives_api_boundary
tests/test_api_cloud_dispatch.py::test_cloud_override_survives_api_boundary_on_local_server
tests/test_api_cloud_dispatch.py::test_cloud_override_survives_api_boundary_on_local_server
tests/test_api_cloud_dispatch.py::test_full_api_process_job_result_carries_cloud_provenance
tests/test_api_cloud_dispatch.py::test_full_api_process_job_result_carries_cloud_provenance
C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\site-
packages\google\auth\_default.py:113: UserWarning: Your application has authenticated using end
user credentials from Google Cloud SDK without a quota project. You might receive a "quota
exceeded" or "API not enabled" error. See the following page for troubleshooting:
https://cloud.google.com/docs/authentication/adc-troubleshooting/user-creds.
warnings.warn(_CLOUD_SDK_CREDENTIALS_WARNING)

```

```

-- Docs: https://docs.pytest.org/en/stable/how-to/capture-warnings.html
===== short test summary info =====
FAILED tests/test_config_deployment.py::test_default_config_dump_only_adds_deployment_section
FAILED tests/test_config_deployment.py::test_save_default_config_includes_deployment
2 failed, 195 passed, 8 warnings in 64.34s (0:01:04)
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

137. assistant (2026-07-08T05:29:48.024Z)

Only 2 failures ? both are config-shape snapshot tests that need the new additive fields. Let me read and update them:

138. user (2026-07-08T05:29:49.424Z)

```

63: def test_default_config_dump_only_adds_deployment_section():
392: def test_save_default_config_includes_deployment(tmp_path):
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

139. user (2026-07-08T05:29:52.089Z)

```

63     def test_default_config_dump_only_adds_deployment_section():
64         """Adding the section must not perturb any other default ? the dump equals a fresh
65         AppConfig()'s dump, and the new section is present with empty values."""
66         dumped = load_config(NONEXISTENT).model_dump(mode="json")
67         assert dumped == AppConfig().model_dump(mode="json")
68         assert dumped["deployment"] == {
69             "profile": "",
70             "compute_target": "",
71             "max_concurrent_remote_jobs": 4,
72             "remote_poll_max_wait_sec": 3900.0,
73         }
74
75
76     # -----
77     # Parity: the model Literal must match the presets table (no drift).
78     # -----
79     def test_profile_literal_matches_presets():
80         import typing
81
82         literal_profiles = set(
83             typing.get_args(DeploymentConfig.model_fields["profile"].annotation)
84         )
85         assert literal_profiles == {"", *PROFILES}
86         assert set(PROFILE_PRESETS) == set(PROFILES)
87         assert set(COMPUTE_TARGET_FOR_PROFILE) == set(PROFILES)
88
89
90     def test_compute_target_literal_matches_presets():
91         import typing

```

```

92
93     literal_targets = set(
94         typing.get_args(DeploymentConfig.model_fields["compute_target"].annotation)
95     )
96     assert literal_targets == {"", *COMPUTE_TARGETS}
97     assert set(COMPUTE_TARGET_FOR_PROFILE.values()) == set(COMPUTE_TARGETS)
98
99
100     def test_each_preset_only_touches_existing_knobs():
101         """A preset must only set knobs that exist on the typed models ? otherwise it would
be a
102         silent no-op. Validates the merged result against AppConfig for every profile."""

```

140. user (2026-07-08T05:30:01.916Z)

The file C:\kg-wt521\tests\test_config_deployment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

141. user (2026-07-08T05:30:04.968Z)

```

395     def test_save_default_config_includes_deployment(tmp_path):
396         p = save_default_config(str(tmp_path / "out.json"))
397         saved = json.loads(p.read_text())
398         assert saved["deployment"] == {
399             "profile": "",
400             "compute_target": "",
401             "max_concurrent_remote_jobs": 4,
402             "remote_poll_max_wait_sec": 3900.0,
403         }
404         cfg = load_config(str(p))
405         assert cfg.deployment.profile == ""
406
407
408     def test_repo_config_json_has_no_profile():
409         """The shipped config.json must not silently opt into a profile (default path stays
today's
410         behaviour for the real deployment)."""
411         repo_root = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
412         cfg_path = os.path.join(repo_root, "config.json")
413         if not os.path.exists(cfg_path):
414             pytest.skip("no repo config.json")
415         cfg = load_config(cfg_path)
416         assert cfg.deployment.profile == ""
417
418
419     def test_remote_poll_max_wait_default_and_validation():
420         """CP-rebuild §4.3: the dispatch poll budget must outlast the worker Job's
421         --task-timeout (3600s) by default, and a non-positive budget is a config error."""
422         from pydantic import ValidationError
423
424         from backend.config.models import DeploymentConfig

```

142. user (2026-07-08T05:30:10.312Z)

The file C:\kg-wt521\tests\test_config_deployment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

143. assistant (2026-07-08T05:30:16.989Z)

Config-shape tests fixed. Now let me write the new tests. First, the phase-placement resolver + guard tests:

144. assistant (2026-07-08T05:30:38.981Z)

```
"""Config-driven phase ? machine placement (#521).
```

```
Proves the ``deployment.phase_placement`` knob + ``resolve_phase_placement`` resolver: the
legacy
defaults reproduce today's placement exactly (validation + compile on the control-plane, segment
on
the L4 iff compute_target==cloud_run), an explicit knob overrides, the cloud-serving preset
places
compile on the L4, and the "nobody validates" footgun is rejected at config load.
"""
```

```
import pytest
```

```
from backend.config import (
    PHASE_COMPILE,
    PHASE_SEGMENT,
    PHASE_VALIDATION,
    PLACE_CLOUD_L4,
    PLACE_CONTROL_PLANE,
    AppConfig,
    load_config,
    resolve_phase_placement,
)
```

----- legacy defaults (byte-identical)

```
def test_default_placement_is_todays_behaviour():
    """No phase_placement block ? validation + compile on the control-plane, segment on the
    control-plane (compute_target unset). Byte-identical to pre-#521."""
    cfg = AppConfig()
    assert resolve_phase_placement(cfg, PHASE_VALIDATION) == PLACE_CONTROL_PLANE
    assert resolve_phase_placement(cfg, PHASE_SEGMENT) == PLACE_CONTROL_PLANE
    assert resolve_phase_placement(cfg, PHASE_COMPILE) == PLACE_CONTROL_PLANE

def test_segment_default_follows_compute_target():
    """The legacy segment default is cloud IFF compute_target==cloud_run ? so a cloud_run server
    keeps dispatching DETECT without any phase_placement config."""
    cloud = AppConfig.model_validate({"deployment": {"compute_target": "cloud_run"}})
    assert resolve_phase_placement(cloud, PHASE_SEGMENT) == PLACE_CLOUD_L4
    # validation + compile still default to the control-plane on a cloud_run server (unset
    knobs).
    assert resolve_phase_placement(cloud, PHASE_VALIDATION) == PLACE_CONTROL_PLANE
    assert resolve_phase_placement(cloud, PHASE_COMPILE) == PLACE_CONTROL_PLANE

    local = AppConfig.model_validate({"deployment": {"compute_target": "local_gpu"}})
    assert resolve_phase_placement(local, PHASE_SEGMENT) == PLACE_CONTROL_PLANE
```

----- explicit overrides

```
def test_explicit_knob_overrides_default():
    """Each phase's explicit knob wins over the legacy default."""
    cfg = AppConfig.model_validate(
        {
            "deployment": {
                "compute_target": "cloud_run",
                "phase_placement": {
                    "segment": "control_plane", # force DETECT back on the control-plane
                    "compile": "cloud_run_l4", # move stitch to the L4
                }
            }
        })
```

```

        },
    }
}
)
assert resolve_phase_placement(cfg, PHASE_SEGMENT) == PLACE_CONTROL_PLANE
assert resolve_phase_placement(cfg, PHASE_COMPILE) == PLACE_CLOUD_L4

def test_resolver_never_raises_on_missing_deployment():
    """A config-ish object with no deployment/phase_placement degrades to the legacy default."""

    class _Bare:
        pass

    assert resolve_phase_placement(_Bare(), PHASE_COMPILE) == PLACE_CONTROL_PLANE
    assert resolve_phase_placement(_Bare(), PHASE_SEGMENT) == PLACE_CONTROL_PLANE

```

----- the cloud-serving preset

```

def test_cloud_serving_preset_places_compile_on_l4(tmp_path):
    """The cloud-serving profile preset moves compile to the L4 (the #521 headline) while
    keeping
    validation on the control-plane and segment on the L4 (via compute_target)."""
    cfg_path = tmp_path / "c.json"
    cfg_path.write_text('{"deployment": {"profile": "cloud-serving"}}')
    cfg = load_config(str(cfg_path))
    assert resolve_phase_placement(cfg, PHASE_COMPILE) == PLACE_CLOUD_L4
    assert resolve_phase_placement(cfg, PHASE_SEGMENT) == PLACE_CLOUD_L4
    assert resolve_phase_placement(cfg, PHASE_VALIDATION) == PLACE_CONTROL_PLANE

```

----- the "nobody validates" guard

```

def test_validation_on_l4_requires_cloud_segment():
    """validation=cloud_run_l4 with segment on the control-plane would leave NOBODY validating
    (the
    control-plane skips + no worker runs) ? rejected at config load."""
    with pytest.raises(Exception, match="phase_placement.validation"):
        AppConfig.model_validate(
            {
                "deployment": {
                    "compute_target": "",
                    "phase_placement": {
                        "validation": "cloud_run_l4",
                        "segment": "control_plane",
                    },
                },
            }
        )

def test_validation_on_l4_ok_with_cloud_segment():
    """validation=cloud_run_l4 is accepted when segment also runs on the L4 (explicit)."""
    cfg = AppConfig.model_validate(
        {
            "deployment": {
                "phase_placement": {
                    "validation": "cloud_run_l4",
                    "segment": "cloud_run_l4",
                }
            }
        }
    )

```



```

        }
    }
)
assert resolve_phase_placement(cfg, PHASE_VALIDATION) == PLACE_CLOUD_L4

def test_validation_on_l4_ok_when_compute_target_is_cloud_run():
    """The guard honours the LEGACY segment default: with compute_target=cloud_run the segment
runs
on the L4 by default, so validation=cloud_run_l4 needs no explicit segment knob."""
    cfg = AppConfig.model_validate(
        {
            "deployment": {
                "compute_target": "cloud_run",
                "phase_placement": {"validation": "cloud_run_l4"},
            }
        }
    )
    assert resolve_phase_placement(cfg, PHASE_VALIDATION) == PLACE_CLOUD_L4
    assert resolve_phase_placement(cfg, PHASE_SEGMENT) == PLACE_CLOUD_L4

```

145. user (2026-07-08T05:30:39.022Z)

File created successfully at: C:\kg-wt521\tests\test_phase_placement.py (file state is current in your context ? no need to Read it back)

146. assistant (2026-07-08T05:31:07.339Z)

Now the compile-dispatch test mirroring `test\_api\_cloud\_dispatch.py`:

147. assistant (2026-07-08T05:31:44.836Z)

"""Compile/stitch ? L4 worker dispatch (#521).

Proves the control-plane can OFFLOAD the reel compile/stitch to the Cloud Run L4 worker by config (`deployment.phase\_placement.compile=cloud\_run\_l4`), mirroring the DETECT dispatch ? all GPU-free + cloud-free via a FAKE `CloudRunJobClient` that simulates the compile worker (reads the staged compile-spec + writes the reel + done-marker into a local "bucket"). Cardinal invariant: DEFAULT (compile on the control-plane) still stitches in-process, byte-identically.

"""

```

import json
import os

```

```

import pytest

```

```

import backend.compute as compute_pkg
import backend.orchestration as orch
from backend.compute.cloud_run import CloudRunJobClient
from backend.config import AppConfig
from backend.infrastructure.database import Database
from backend.infrastructure.execution_policy import ExecutionPolicy

```

```

_FAST = ExecutionPolicy(poll_every_n_sec=0.0, op_timeout_sec=5.0, max_wait_sec=5.0)

```

```

class _FakeCompileClient(CloudRunJobClient):
    """Simulates the Cloud Run compile worker: on dispatch, READ the staged compile-spec
(proving the
control-plane resolved + staged it), then write the reel done-marker into the local bucket
exactly
where the real worker would, so the dispatcher's bucket-poll resolves."""

```

```

def __init__(self, *, returncode=0, video_id=None):
    self.calls: list[list[str]] = []
    self.spec: dict | None = None
    self.returncode = returncode
    self._video_id = video_id

def run_job(self, job_name, argv, *, region, project=None):
    self.calls.append(list(argv))
    spec_uri = argv[argv.index("--spec-uri") + 1]
    done_uri = argv[argv.index("--done-uri") + 1]
    with open(spec_uri, encoding="utf-8") as f:
        self.spec = json.load(f)
    vid = self._video_id or self.spec["video_id"]
    out_name = self.spec["output_name"]
    status = "success" if self.returncode == 0 else "failed"
    marker = {
        "video_id": vid,
        "returncode": self.returncode,
        "status": status,
        "download_url": f"/api/highlights/{vid}/{out_name}",
        "reel_uri": f"gs://bucket/highlights/{vid}/{out_name}",
        "message": "" if self.returncode == 0 else "stitch blew up",
    }
    os.makedirs(os.path.dirname(done_uri), exist_ok=True)
    with open(done_uri, "w", encoding="utf-8") as f:
        json.dump(marker, f)
    return "exec-compile-1"

def cancel_execution(self, execution): # pragma: no cover - happy path never cancels
    raise AssertionError("cancel must never fire on the happy path")

def _inject_fake(monkeypatch, fake):
    real = compute_pkg.get_compute_backend
    monkeypatch.setattr(
        compute_pkg,
        "get_compute_backend",
        lambda config, **kw: real(config, run_client=fake, policy=_FAST, token="tok", **kw),
    )

@pytest.fixture
def compile_env(tmp_path, monkeypatch):
    """A control-plane DB with a video + 3 play segments, wired for cloud with compile placed on
    the
    L4 and a LOCAL bucket (so the fake worker's spec/marker are plain file reads/writes)."""
    db = Database(str(tmp_path / "t.db"))
    bucket = tmp_path / "bucket"
    bucket.mkdir()
    video_id = "vidCompile"
    src = tmp_path / "source.mp4"
    src.write_bytes(b"SRC")
    db.add_video(video_id, str(src), "processed", [])
    sid1 = db.add_play_segment(video_id, 2.0, 6.0, metadata={"shot_count": 5})
    sid2 = db.add_play_segment(video_id, 10.0, 14.0, metadata={"shot_count": 3})
    sid3 = db.add_play_segment(video_id, 20.0, 22.0, metadata={"shot_count": 4})
    cfg = AppConfig.model_validate(
        {
            "deployment": {
                "compute_target": "cloud_run",
                "phase_placement": {"compile": "cloud_run_l4"},
            },
            "storage": {"backend": "local", "output_root": str(bucket)},
        }
    )

```

```

return db, cfg, video_id, [sid1, sid2, sid3], bucket, monkeypatch

def _job(video_id, segment_ids, name="reel.mp4", locale=None):
    return {
        "id": "cj1",
        "video_id": video_id,
        "kind": "compile",
        "params": json.dumps(
            {"segment_ids": segment_ids, "output_filename": name, "locale": locale}
        ),
    }

def test_compile_dispatches_to_l4_and_returns_reel(compile_env):
    db, cfg, video_id, seg_ids, _bucket, monkeypatch = compile_env
    fake = _FakeCompileClient()
    _inject_fake(monkeypatch, fake)

    out = orch._execute_compile(cfg, db, _job(video_id, seg_ids, "myreel.mp4"))

    assert out["status"] == "success"
    assert out["video_id"] == video_id
    assert out["download_url"] == f"/api/highlights/{video_id}/myreel.mp4"
    assert out["compute"]["path"] == "cloud_run"
    assert out["compute"]["execution"] == "exec-compile-1"
    assert len(fake.calls) == 1 # exactly one Cloud Run Job execution

def test_compile_dispatch_stages_resolved_spec_with_nvenc(compile_env):
    """The staged compile-spec carries the RESOLVED time-pairs (padding-agnostic start/end from
    the DB)
    and the L4 encoder (deployment.cloud_stitch_encoder = h264_nvenc by default) ? so the worker
    hardware-encodes and needs no control-plane DB."""
    db, cfg, video_id, seg_ids, _bucket, monkeypatch = compile_env
    fake = _FakeCompileClient()
    _inject_fake(monkeypatch, fake)

    orch._execute_compile(cfg, db, _job(video_id, seg_ids))

    assert fake.spec is not None
    assert fake.spec["video_id"] == video_id
    assert fake.spec["video_uri"].endswith("source.mp4")
    # 3 resolved (start, end) pairs from the DB rows
    assert fake.spec["time_pairs"] == [[2.0, 6.0], [10.0, 14.0], [20.0, 22.0]]
    # the L4 stitch uses NVENC (forwarded from deployment.cloud_stitch_encoder)
    assert fake.spec["stitching"]["encoder"] == "h264_nvenc"
    # the argv is a `compile` subcommand carrying the spec + done marker
    argv = fake.calls[0]
    assert argv[0] == "compile"
    assert "--spec-uri" in argv and "--done-uri" in argv
    joined = " ".join(argv)
    assert "storage.output_root=" in joined # bucket forwarded so the worker mirrors the reel

def test_compile_dispatch_forwards_stitch_config_overrides(compile_env):
    """The worker Job has no DEPLOYMENT_PROFILE, so the dispatch forwards the bucket coordinates
    (storage.output_root/backend) ? else the worker can't read the spec or mirror the reel."""
    db, cfg, video_id, seg_ids, bucket, monkeypatch = compile_env
    fake = _FakeCompileClient()
    _inject_fake(monkeypatch, fake)
    orch._execute_compile(cfg, db, _job(video_id, seg_ids))
    joined = " ".join(fake.calls[0])
    assert f"storage.output_root={bucket}" in joined
    assert "storage.backend=local" in joined

```

```

def test_compile_cloud_failure_returns_failed(compile_env):
    """A worker stitch failure (rc != 0) surfaces as a failed job with the worker's message ?
the SPA
    gets a clean verdict, not a crash."""
    db, cfg, video_id, seg_ids, _bucket, monkeypatch = compile_env
    fake = _FakeCompileClient(returncode=3)
    _inject_fake(monkeypatch, fake)

    out = orch._execute_compile(cfg, db, _job(video_id, seg_ids))
    assert out["status"] == "failed"
    assert "stitch blew up" in out["message"]
    assert out["video_id"] == video_id

def test_compile_default_stitches_in_process_no_dispatch(compile_env, monkeypatch):
    """DEFAULT (compile on the control-plane): _execute_compile stitches in-process via
    _stitch_reel
    and NEVER dispatches to Cloud Run. The stitch core is stubbed so no real ffmpeg runs."""
    db, cfg, video_id, seg_ids, _bucket, _mp = compile_env
    # flip compile back to the control-plane (the default) ? everything else stays cloud-wired
    cfg.deployment.phase_placement.compile = "control_plane"
    fake = _FakeCompileClient()
    _inject_fake(monkeypatch, fake)
    seen = {}

    def _fake_stitch(c, stitching, vid, source, pairs, name, locale, notify):
        seen["pairs"] = pairs
        seen["source"] = source
        return {"status": "success", "video_id": vid, "download_url":
f"/api/highlights/{vid}/{name}"}

    monkeypatch.setattr(orch, "_stitch_reel", _fake_stitch)

    out = orch._execute_compile(cfg, db, _job(video_id, seg_ids, "local.mp4"))
    assert out["status"] == "success"
    assert fake.calls == [] # NEVER dispatched ? the control-plane stitched in-process
    assert seen["pairs"] == [(2.0, 6.0), (10.0, 14.0), (20.0, 22.0)]
    assert seen["source"].endswith("source.mp4")

def test_compile_missing_video_fails_before_dispatch(compile_env):
    """A compile for an unknown video_id fails at resolve (404) before any dispatch."""
    db, cfg, _video_id, _seg_ids, _bucket, monkeypatch = compile_env
    fake = _FakeCompileClient()
    _inject_fake(monkeypatch, fake)
    out = orch._execute_compile(cfg, db, _job("nope", [1, 2]))
    assert out["status"] == "failed"
    assert fake.calls == [] # never dispatched ? resolve failed first

```

148. user (2026-07-08T05:31:44.880Z)

File created successfully at: C:\kg-wt521\tests\test_compile_dispatch.py (file state is current in your context ? no need to Read it back)

149. assistant (2026-07-08T05:31:50.238Z)

Let me run the new placement and compile-dispatch tests:

150. user (2026-07-08T05:31:57.622Z)

```

)

# Build the compile-spec the worker consumes. time-pairs are JSON lists (frozen-safe on
the wire);

```

```

# the stitching config is the control-plane's RESOLVED one, with the encoder swapped to
the L4's
# (deployment.cloud_stitch_encoder, NVENC by default) so the worker hardware-encodes the
reel.
stitch_dump = stitching.model_dump()
stitch_dump["encoder"] = getattr(
    getattr(cfg, "deployment", None), "cloud_stitch_encoder", "h264_nvenc"
)
spec_payload = {
    "video_id": video_id,
    "video_uri": video["filepath"],
    "output_name": out_name,
    "locale": locale,
    "time_pairs": [
        [float(s["start_time"]), float(s["end_time"])] for s in kept
    ],
    "stitching": stitch_dump,
}
> token = uuid.uuid4().hex
      ^^^^
E      NameError: name 'uuid' is not defined. Did you forget to import 'uuid'?

```

```

backend\orchestration.py:1339: NameError
===== short test summary info =====
FAILED tests/test_compile_dispatch.py::test_compile_dispatches_to_l4_and_returns_reel
FAILED tests/test_compile_dispatch.py::test_compile_dispatch_stages_resolved_spec_with_nvenc
FAILED tests/test_compile_dispatch.py::test_compile_dispatch_forwards_stitch_config_overrides
FAILED tests/test_compile_dispatch.py::test_compile_cloud_failure_returns_failed
4 failed, 10 passed in 3.28s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

151. assistant (2026-07-08T05:32:00.933Z)

`uuid` isn't imported in orchestration.py. Let me add it:

152. user (2026-07-08T05:32:02.298Z)

The file C:\kg-wt521\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

153. user (2026-07-08T05:32:10.764Z)

```

..... [100%]
14 passed in 2.88s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

154. assistant (2026-07-08T05:32:20.686Z)

All 14 pass. Now add `run\_compile` tests to `test\_compute\_backend.py`. Let me register the worktree file's tail:

155. user (2026-07-08T05:32:21.656Z)

```

1094     def heartbeat(_elapsed):
1095         pass
1096
1097     result = backend.run_infer(
1098         ComputeJobSpec(video_uri="gs://b/v.mp4", out_uri=str(tmp_path / "bk")),
1099         on_wait=heartbeat,
1100     )
1101     assert seen["on_tick"] is heartbeat
1102     assert result.video_id == "vidCloud"
1103

```

156. user (2026-07-08T05:32:44.926Z)

The file C:\kg-wt521\tests\test_compute_backend.py has been updated successfully. (file state is current in your context ? no need to Read it back)

157. assistant (2026-07-08T05:32:48.282Z)

Let me remove the unused import in that last test:

158. user (2026-07-08T05:32:49.760Z)

The file C:\kg-wt521\tests\test_compute_backend.py has been updated successfully. (file state is current in your context ? no need to Read it back)

159. assistant (2026-07-08T05:32:55.406Z)

Now the worker `cmd\_compile` + parser tests. Let me find the end of `test\_worker.py`:

160. user (2026-07-08T05:32:57.145Z)

```
945 tests/test_worker.py
    ) # graceful: _resolve_train_detector caught NotImplementedError, returned None

def test_probe_video_fps_width_returns_none_on_unreadable(tmp_path):
    """The strict ffprobe probe must return None (caller skips) for a non-video file, never
    guess."""
    bad = tmp_path / "not_a_video.mp4"
    bad.write_bytes(b"NOTAVIDEO")
    assert wmain._probe_video_fps_width(str(bad)) is None
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

161. user (2026-07-08T05:33:00.234Z)

```
938         ) # graceful: _resolve_train_detector caught NotImplementedError, returned None
939
940
941     def test_probe_video_fps_width_returns_none_on_unreadable(tmp_path):
942         """The strict ffprobe probe must return None (caller skips) for a non-video file,
943         never guess."""
944         bad = tmp_path / "not_a_video.mp4"
945         bad.write_bytes(b"NOTAVIDEO")
946         assert wmain._probe_video_fps_width(str(bad)) is None
```

162. user (2026-07-08T05:33:27.700Z)

The file C:\kg-wt521\tests\test_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

163. assistant (2026-07-08T05:33:31.875Z)

test_worker.py needs `import os`. Let me add it:

164. user (2026-07-08T05:33:32.842Z)

The file C:\kg-wt521\tests\test_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

165. assistant (2026-07-08T05:33:35.717Z)

Now add an encoder-selection test to `test\_compiler.py`:

166. user (2026-07-08T05:33:37.489Z)

```
def test_compile_raises_when_source_missing(tmp_path):
    comp = VideoCompiler(str(tmp_path / "out"))
    with pytest.raises(FileNotFoundError):
        comp.compile_highlights(str(tmp_path / "nope.mp4"), [(0.0, 5.0)], "reel.mp4")
---lines---
160 tests/test_compiler.py
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

167. user (2026-07-08T05:33:42.130Z)

```
155
156
157     def test_compile_raises_when_source_missing(tmp_path):
158         comp = VideoCompiler(str(tmp_path / "out"))
159         with pytest.raises(FileNotFoundError):
160             comp.compile_highlights(str(tmp_path / "nope.mp4"), [(0.0, 5.0)], "reel.mp4")
```

168. user (2026-07-08T05:34:00.347Z)

The file C:\kg-wt521\tests\test_compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

169. assistant (2026-07-08T05:34:03.208Z)

Let me run all the new/modified test files:

170. user (2026-07-08T05:34:15.021Z)

```
..... [ 53%]
..... [100%]
135 passed in 8.14s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

171. assistant (2026-07-08T05:34:47.831Z)

All 135 pass. Now let me add tests for the validation-placement wiring (the `skip\_validation` flip + CP delegation) to `test\_api\_cloud\_dispatch.py`:

172. user (2026-07-08T05:34:48.712Z)

```
1097     def test_wait_progress_message_is_honest_and_minute_granular():
1098         """#482: the CP knows ELAPSED time, not frames ? the heartbeat says so plainly,
shows
1099     movement from the first tick ("starting up" during queue/cold-start), then counts
minutes."""
1100     from backend.orchestration import _wait_progress_message
1101
1102     early = _wait_progress_message("wasb_hybrid", 12.0)
1103     assert "wasb_hybrid" in early and "starting up" in early
1104     later = _wait_progress_message("wasb_hybrid", 125.0)
1105     assert later.endswith("2 min elapsed")
1106
```

173. user (2026-07-08T05:35:03.239Z)

The file C:\kg-wt521\tests\test_api_cloud_dispatch.py has been updated successfully. (file state is current in your context ? no need to Read it back)

174. user (2026-07-08T05:35:41.682Z)

```
tests/test_api_cloud_dispatch.py::test_full_api_process_to_query_via_cloud
tests/test_api_cloud_dispatch.py::test_local_override_survives_api_boundary
tests/test_api_cloud_dispatch.py::test_local_override_survives_api_boundary
```

```
tests/test_api_cloud_dispatch.py::test_cloud_override_survives_api_boundary_on_local_server
tests/test_api_cloud_dispatch.py::test_cloud_override_survives_api_boundary_on_local_server
tests/test_api_cloud_dispatch.py::test_full_api_process_job_result_carries_cloud_provenance
tests/test_api_cloud_dispatch.py::test_full_api_process_job_result_carries_cloud_provenance
C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\site-
packages\google\auth\_default.py:113: UserWarning: Your application has authenticated using end
user credentials from Google Cloud SDK without a quota project. You might receive a "quota
exceeded" or "API not enabled" error. See the following page for troubleshooting:
https://cloud.google.com/docs/authentication/adc-troubleshooting/user-creds.
warnings.warn(_CLOUD_SDK_CREDENTIALS_WARNING)
```

```
-- Docs: https://docs.pytest.org/en/stable/how-to/capture-warnings.html
40 passed, 8 warnings in 33.32s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

175. assistant (2026-07-08T05:35:49.665Z)

All 40 pass. Now let me run lint + type checks (matching CI). Let me check the exact CI commands and run them:

176. user (2026-07-08T05:35:51.173Z)

```
16: # Lint gate (ruff.toml at repo root): pyflakes/pycodestyle + bugbear.
17: # Also the repo-wide syntax guard ? ruff reports invalid-syntax on every
29:     - name: Install ruff
30:         # Pinned so an upstream ruff release can't redden every PR overnight; bump
deliberately.
31:         run: python -m pip install ruff==0.15.16
33:         run: ruff check . --output-format github
35: # PII / attribution leak guard (tools/leak_scan.py): fails on a committed 32-hex MD5
49:         run: python -m tools.leak_scan
51: # Internal-link guard (tools/check_md_links.py): fails on a broken internal markdown link
64:         run: python -m tools.check_md_links
66: # Type gate (mypy.ini at repo root): lenient green baseline with an explicit
68: # up. Typed core deps are installed so mypy sees their real signatures
79:     - name: Install mypy + typed core deps
80:         # mypy pinned alongside ruff so a type-checker release can't redden CI overnight.
81:         run: python -m pip install mypy==2.1.0 pydantic fastapi uvicorn python-dotenv google-
genai numpy
83:         run: mypy backend training
136:         # (pytest, pytest-cov, httpx for the fastapi TestClient, ruff, mypy).
146:         run: python run_all_tests.py --cov=backend --cov=training --cov-report=term --cov-
fail-under=70
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

177. user (2026-07-08T05:35:55.377Z)

```
=== RUFF ===
I001 [*] Import block is un-sorted or un-formatted
--> tests/test_phase_placement.py:9:1
|
7 |     """
8 |
9 | / import pytest
10 | |
11 | | from backend.config import (
12 | |     PHASE_COMPILE,
13 | |     PHASE_SEGMENT,
14 | |     PHASE_VALIDATION,
15 | |     PLACE_CLOUD_L4,
16 | |     PLACE_CONTROL_PLANE,
17 | |     AppConfig,
18 | |     load_config,
19 | |     resolve_phase_placement,
```



```
20 | | )
    | | ^
    |
```

help: Organize imports

Found 1 error.

[*] 1 fixable with the `--fix` option.

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

178. assistant (2026-07-08T05:35:58.315Z)

One import-sort nit. Let me auto-fix and re-check:

179. user (2026-07-08T05:36:00.348Z)

Found 1 error (1 fixed, 0 remaining).

=== RE-CHECK ===

All checks passed!

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

[This command modified 1 file you've previously read: ..\..\..\..\kg-wt521\tests\test_phase_placement.py. Call Read before editing.]

180. assistant (2026-07-08T05:36:03.325Z)

Ruff clean. Now mypy on the changed modules:

181. user (2026-07-08T05:37:09.208Z)

```
backend\annotations\base.py:88: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\sports\__init__.py:9: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\pipeline\stitcher\compiler.py:57: error: Value of type variable
"SupportsRichComparisonT" of "min" cannot be "float | None" [type-var]
backend\pipeline\validators\base.py:123: note: By default the bodies of untyped functions are
not checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Found 1 error in 1 file (checked 150 source files)
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

182. assistant (2026-07-08T05:37:17.871Z)

mypy flags the probe's `min(30, ffmpeg\_timeout\_sec())` (the timeout can be `None`). Let me fix it:

183. user (2026-07-08T05:37:24.270Z)

The file C:\kg-wt521\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

184. user (2026-07-08T05:37:35.411Z)

=== exit summary ===

```
backend\providers\__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 150 source files
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

185. assistant (2026-07-08T05:37:43.579Z)

mypy clean. Now let me run the FULL test suite with coverage (the pre-LGTM gate), running in the background isn't needed ? let me run it directly:

186. user (2026-07-08T05:41:08.359Z)

backend\storage\gdrive_api.py	105	6	94%
backend\storage\local.py	49	3	94%
backend\storage\ref.py	49	0	100%
backend\storage\s3.py	47	5	89%
backend\tools__init__.py	0	0	100%
backend\tools\export.py	75	2	97%
backend\tools\rally_slicer.py	105	47	55%
backend\utils__init__.py	0	0	100%
backend\utils\annotation_proxy.py	64	11	83%
backend\utils\app_home.py	40	0	100%
backend\utils\ffmpeg.py	47	3	94%
backend\utils\freespace.py	32	2	94%
backend\utils\geometry.py	29	0	100%
backend\utils\hashing.py	9	0	100%
backend\utils\logging_setup.py	11	0	100%
backend\utils\paths.py	44	3	93%
backend\utils\redact.py	16	0	100%
backend\utils\run_telemetry.py	126	12	90%
backend\utils\single_instance.py	43	11	74%
backend\utils\video.py	34	3	91%
backend\utils\workspace.py	49	0	100%
backend\worker__init__.py	0	0	100%
backend\worker__main__.py	371	32	91%
training__init__.py	0	0	100%
training\gen0__init__.py	0	0	100%
training\gen0\harness.py	130	9	93%

TOTAL	12663	1810	86%

Required test coverage of 70% reached. Total coverage: 85.71%

===== short test summary info =====
FAILED tests/test_main_module_budget.py::test_main_py_line_budget - Assertion...
1 failed, 1874 passed, 8 skipped, 25 warnings in 199.57s (0:03:19)

=====

```
RESULT: FAIL - pytest exit 1
```

=====

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

187. assistant (2026-07-08T05:41:17.428Z)

Coverage 85.7% (??0). One failure: the `main.py` line-budget guard (I added 2 re-export lines).
Let me check it:

188. user (2026-07-08T05:41:19.086Z)

```
"""Anti-bloat ratchet for the ``backend.main`` god-module split (P23; re-cut of #390).
```

```
``backend/main.py`` grew 1448 -> 2360 lines in under a week because every endpoint/fix PR  
landed in it. The P23 extraction moved the processing/compile orchestration to  
``backend/orchestration.py``; this ratchet keeps the split from silently regressing:
```

- **Line budget:** ``backend/main.py`` must stay at/below ``MAIN\_LINE\_BUDGET``. When a later split lands (CLI -> P23 step 2, routes -> P24), **TIGHTEN** the budget to the new size + small slack ? never loosen it without an owner-visible PR discussion.
- **Route ceiling:** no NEW ``@app.<method>`` routes in ``main.py`` ? new endpoints go in ``backend/api/`` (see ``backend/api/uploads.py`` for the pattern).
- **Re-export contract:** every orchestration name stays addressable on ``backend.main`` (```import backend.main as m; m.execute_processing``` and ```monkeypatch.setattr(backend.main, ...)``` are load-bearing seams ? see ```backend/api/job_worker.py```'s module docstring).

If this test fails your PR: extract, don't inline. Umbrella tracker (archived):

```
``docs/archives/past_projects/AUDIT_REPORT_khelsutra-guru_2026-07-02-COMPLETED-2026-07-02.md``  
(Remediation tracker P23/P24).
```

```
"""
```

```
import ast
from pathlib import Path
```

```
REPO = Path(__file__).resolve().parents[1]
MAIN = REPO / "backend" / "main.py"
```

main.py was 1623 lines after the P24 server_runtime extraction (2026-07-02; 1703 after P23 orchestration extraction). Slack covers small in-place fixes; systematic growth moved a module instead (orchestration -> backend/orchestration.py; server/launcher -> backend/launcher.py; new routes -> backend/api/). TIGHTEN this budget whenever a further split lands.

1650 -> 1662 (#448): the /api/jobs list ROUTE was extracted to backend/api/jobs.py (route budget respected); the only in-place growth was the idempotent-submit dedup hook in process.py.

1662 -> 1653 (#481 review): the dedup became an ATOMIC single-op claim (backend/api/jobs.py + db.claim_or_create_process_job), which REPLACED the inline create_job block ? reworked.

```
MAIN_LINE_BUDGET = 1653
```

The 18 routes present at the P23 extraction. New endpoints belong in backend/api/.

```
MAIN_ROUTE_BUDGET = 18
```

```
ORCHESTRATION_REEXPORTS = [
    "_cloud_available",
    "_CompileError",
    "_execute_compile",
    "_execute_processing",
    "_finalize_reingest",
    "_ingest_remote_segments",
    "_maybe_dispatch_remote",
    "_remote_input_size_bytes",
    "_report_config_for_input",
    "_resolve_compile",
]
```

```
def test_main_py_line_budget():
    n = sum(1 for _ in MAIN.open(encoding="utf-8"))
    assert n <= MAIN_LINE_BUDGET, (
        f"backend/main.py is {n} lines (budget {MAIN_LINE_BUDGET}). Don't grow the "
        f"god-module back ? extract to backend/orchestration.py / backend/api/ / a new "
        f"module instead (P23/P24 in docs/archives/past_projects/"
    )
    Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

189. assistant (2026-07-08T05:41:32.967Z)

The two re-exports I added aren't load-bearing (nothing calls `m.\_stitch\_reel`; the worker uses `orchestration.\_stitch\_reel` directly), and they push `main.py` over its anti-bloat budget. Let me revert them:

190. user (2026-07-08T05:41:37.300Z)

The file C:\kg-wt521\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

191. user (2026-07-08T05:41:51.985Z)

```
...
3 passed in 2.67s
```

[100%]

```
=== leak_scan ===
[leak-scan] clean ? no MD5 / personal-path / private-name leaks in the code surface.
=== md_links ===
[check-md-links] clean ? no broken internal links in live docs.
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

192. assistant (2026-07-08T05:41:56.130Z)

All gates green. Now the docs. Let me read the COMPUTE_DECOUPLED_SERVING README to place the new doc:

193. user (2026-07-08T05:41:57.794Z)

Compute-Decoupled Rally Serving ? one pure core, profile-selected compute + storage

```
> **Status:** ? **M0 APPROVED** (owner, 2026-06-21) ? **IN PROGRESS (build M1+)** . **Owner:**
Avi . **Created:** 2026-06-21 . **Last updated:** 2026-06-21
> **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/past_projects/`
when DONE)
> **Tracking anchors:** §7 progress tracker is the source of truth; pointer in memory `current-
state` + `docs/NEXT_STEPS.md` once work starts.
> **Relation to existing docs:** the productionization successor to the archived
`DECOUPLED_COMPUTE` (M3/M4); realizes the `ComputeBackend`/`StorageBackend` registries from
[PLATFORM_ARCHITECTURE.md](../PLATFORM_ARCHITECTURE.md) §3; the truly-local profile is the
L0 tier of [VIDEO_LOCALITY_MODEL.md](../VIDEO_LOCALITY_MODEL.md) and the `LocalDesktop` backend
of [DESKTOP_APP_PLAN.md](../DESKTOP_APP_PLAN.md).
> **Naming:** this was the "Cloud Run + GPU serving subproject" (ADR-012). Renamed compute-
decoupled serving because truly-local is a co-equal, first-class profile, not an
afterthought.
> **Honesty note:** claims tagged `[verified]` (checked against code by the design workflow
`wf_56274ff0`), `[design]` (proposed), `[researched]`.

---
```

```
> Northstar (D14): a mobile, app-like flow ? point at a Google Drive video ?
fully cloud-native run ? the stitched, ready-to-share reel lands back in Drive, next
to the source. Two design promises around it: a user can switch local?cloud anytime (D13,
a headline feature), and execution is never a surprise ? auto-detect suggests, the user
confirms (D15).
```

0. TL;DR

The rally engine is a **pure function of** `(input bytes + config) ? outputs` in three deterministic stages ? **DETECT ? WINDOW ? STITCH** ? that **no deployment profile may branch on**. A **ComputeBackend** (deployment profile) decides **where** the core runs (a local process vs a Cloud Run GPU job); a **StorageBackend** decides **where bytes come from** (local FS / USB / mounted Google Drive / bucket). **One startup config** picks both. We ship **truly-local** (local GPU + in-process stitch, zero cloud) and **cloud-serving** (upload `<2GB` / gdrive / bucket ? a Cloud Run GPU job runs detect **and** stitch, so stitch's compute is metered), plus **cpu-mock** for CI. The single most important caveat: detection is **byte-identical only** on the same GPU (cross-GPU is sub-pixel-different, per the 2026-06-20 deploy report) ? so the **parity guarantee** is enforced at the WINDOW + STITCH layer, not detector-CSV bytes.

1. Problem & goal

- **Why now:** DECOUPLED_COMPUTE proved compute **portability** (M0?M2 + M3.5 on a GCP L4) and is archived; its deferred **M3** (serving endpoint + per-video billing) + M4 (scale-to-zero + cost guard) become this project. The owner needs **both** a hosted service **and** a truly-local mode, with the **core** (detect+stitch) unaffected by which one runs.
- **The two user-facing outcomes** (owner, 2026-06-21):
 - **(a) Hosted service** ? operate on small **local uploads** (`<2GB`) **and** load from **gdrive** / a bucket, spinning up **Cloud Run** for both rally segmentation AND stitching (stitching is compute-intensive ? run it where its cost is accounted for).
 - **(b) Truly-local** ? on a local machine with videos on **local drives/USB** + a local GPU, run inference **and** stitching **completely locally**.

- **"Good"** looks like: the same `video ? rallies ? reel` core gives identical windows + stitched ranges whether it runs on a laptop or Cloud Run; switching is a **config/startup choice**, never a code branch in the core.

2. Decisions locked

#	Decision	Source	Implication
D1	The core is a pure 3-stage function (DETECT/WINDOW/STITCH); no profile branch inside it. design <code>[verified]</code> All profile difference lives in config + the two registries.		
D2	Two registries : <code>ComputeBackend</code> (where) selected by <code>deployment.profile</code> + <code>compute_target</code> ; <code>StorageBackend</code> (what-bytes) by <code>input_backend/storage.backend</code> . ADR-014 Realizes PLATFORM_ARCHITECTURE's <code>local/hosted/byo_worker</code> .		
D3	Profiles shipped : <code>truly-local</code> , <code>cloud-serving</code> , <code>cpu-mock</code> (CI); <code>byo_worker</code> reserved/deferred. ADR-014 Enum/seam reserved so BYO isn't precluded.		
D4	In cloud-serving , STITCH runs on Cloud Run (co-located with detect) so CPU/wall + egress are metered into the per-job cost ledger. owner (a) Stitch-on-laptop would hide real cost.		
D5	Parity enforced at WINDOW+STITCH (byte-exact for a fixed trajectory); detect is byte-exact only same-GPU. deploy report 2026-06-20 Cross-GPU asserted at the rally/window level, never CSV bytes.		
D6	Default-OFF , smallest-safe-first, one PR per milestone; any core-caller change (e.g. configurable output base) ships behind a flag + a Launch Report. <code>[[launch-report-convention]]</code> Zero-regression discipline.		
D7	Cloud Run cold-start: SCALE-T0-ZERO ? accept the ~123 min cold-start (it's an async job: upload?poll?reel, amortized over the multi-minute job). No warm pool (that ~\$500/mo idle would break per-video billing). owner 2026-06-21 M6/M7. Revisit a warm lane only on a UX complaint.		
D8	Pricebook = a versioned DB table; cost computed in USD (matches GCP billing = one source of truth), displayed in INR at a configurable FX rate (Bangalore market); re-versioned when GCP prices change. owner 2026-06-21 M8.		
D9	Edge auth = Cloudflare Access (reuse the site's existing CF). Lock the Cloud Run ingress to the CF proxy + verify the Cf-Access JWT signature (so a direct hit to the Cloud Run URL can't spoof the identity header). owner 2026-06-21 M9. One identity system.		
D10	Free-tier "data-for-reels" trade : free reels in exchange for training rights (uploaded footage + derived trajectories/labels); paid tier = no-train (privacy is the paid feature). Carve-outs : train ONLY on attested ADULT footage (minor footage ? reel yes, training pool NO ? DPDP/GDPR need verifiable parental consent); train on DERIVED data + auto-delete raw ; ToS counsel-reviewed ; live training-on-uploads gated to M9 (public signup). Truly-local needs no DPA. owner 2026-06-21 M9 + D12; ties <code>[[paper-coauthor-aim]]</code> / PERSONALIZATION flywheel.		
D11	Quality floor: Gemini handoff ON for cloud-serving until the owned model clears the #172 ship-gate (e.g. ?0.70 multi-court) ? then flip via a Launch Report ; the owned model runs in shadow/eval meanwhile. (Gemini = serving/inference ONLY, never a training source ? CLAUDE.md guardrail.) owner 2026-06-21 M7 + D12.		
D12	Data-flywheel harness is IN SCOPE (owner add, Q5): capture the (consented, adult) uploaded video + the Gemini reel/rally output ? reliably store in the per-video workspace/bucket ? run shadow evals (owned-vs-Gemini, dual-eval-set) ? promote good ones to the golden TRAINING set (human-reviewed labels) so the owned model climbs toward the D11 floor (then Gemini flips off). owner 2026-06-21 new M11 ; ties D10 (consent) + <code>data_pipeline/</code> + the eval/experiment harness.		
D13	Profile is user-switchable ANYTIME ? a headline FEATURE. The same user runs local today, cloud tomorrow, back to local ? it's just a config/startup choice; the pure core gives equivalent output either way. Advertise it ("your machine *or* our cloud ? your call, switch anytime"). owner 2026-06-21 Honest caveat: equivalent at the rally/window level, not bit-identical across GPUs (D5); a single job runs in one profile, switching is between jobs.		
D14	NORTHSTAR ? operate toward this: a mobile, app-like flow ? point at a Google Drive video ? fully cloud-native run ? the stitched, ready-to-share reel lands back in Drive, next to the source. owner 2026-06-21 Implies a <code>gdrive-api</code> (OAuth) StorageBackend (read source + write reel back) ? distinct from the local mount backend; resolves \$8 #7 , new M12 . The mobile UI is the khelsutra track; the engine must support it (async jobs + status + Drive round-trip / signed URL). Net-new scope (OAuth + Drive write);		

not necessarily v1, but the design must not preclude it (the `StorageBackend` ABC accommodates it). |

| D15 | ****Auto-detect SUGGESTS, the user CONFIRMS ? execution is NEVER a surprise.**** A GPU owner may still choose cloud; ****cloud (= spend) is never entered silently****; the active profile is always explicit + surfaced. | owner 2026-06-21 | Amends §4.3 (auto-detect = a proposed default, not an auto-decision). |

| D16 | ****M6/M7 implementation SHAPE + the cheap §8 picks (owner 2026-06-21):**** (a) ****Cloud Run JOBS**** (not Services-with-GPU) for the detect+stitch unit, ****bucket-poll for completion**** (reuses `JobWaiting`/`claim\_due\_job` verbatim, fits async upload?poll?reel + scale-to-zero D7); (b) ****stitch ENCODE = libx264**** (protects the D5 byte-determinism/parity guarantee + the L4 parity test) ? ****NVDEC for DECODE only**** (a speed lever that never touches output bytes); ****NVENC encode is DEFERRED behind a config knob****, revisit only if M7 cost data justifies relaxing parity (§8 Q8); (c) ****input cap = hard-reject`<2GB`**** for v1 (a config cap so raising to 4GB later is one line; matches the shipped M2 chunked-upload threshold ? §8 Q9). | owner 2026-06-21 | M6/M7. Resolves §8 Q2 (co-located, = D4), Q8, Q9, Q10 (dir name confirmed). |

| D17 | ****Real cloud verification BUDGET = \$100 / ?10k cap (owner 2026-06-21):**** the owner authorized actual GCP spend for due-diligence + production-grade testing/verification of M5?M7 (real L4 run, the M5 runlog, the M7 DEPLOY_REPORT), ****provided the cap is respected**** ? confirm exact cmd + hourly cost before each spend, prefer the smallest viable run + cheapest-adequate GPU, and ****DELETE every resource after**** ([[gcp-vm-always-delete]]). Removes the M6/M7 "spend" gate within the envelope; the calibrated pricebook (M8) + counsel ToS (M9) + OAuth app (M12) stay owner/counsel-side. | owner 2026-06-21 | Unblocks the real M5/M7 runs (not just the code) within the cap. |

3. Foundation ? evolution / ADR lineage ? ***trace the idea end-to-end***

This project is the latest step in a long decision chain; each ADR contributed a piece and a ****subtlety that carries forward****:

| ADR / doc | Contributed | Subtlety carried into this project |

|---|---|---|

| ****ADR-005**** | Permissive licensing (MIT/Apache-2.0/BSD only) | The engine + any served model + the ****container image**** (ffmpeg, fonts, weights) stay commercial-clean. |

| ****ADR-008**** | Owned-model topology; ****train==serve featurizer single-sourcing**** | The "core unaffected across profiles" is the same parity discipline, extended from train/serve to ****local/cloud****. |

| ****ADR-009**** | KhelSutra Guru; ****local-first delivery**** | The ****truly-local profile**** is the direct continuation ? privacy/offline self-host stays first-class. |

| ****ADR-010**** | DECOUPLED: D0?GCP pivot | The cloud-compute origin (GPU + co-located bucket). |

| ****ADR-011**** | DECOUPLED scope = M0?M2+M3.5; ****deferred M3 (serving) + M4 (billing/cost-guard)****; overrode the ****"rent nothing / not a service" posture** | ****This project realizes M3/M4.**** Carries §06's owner-gated gates: ****DPA/consent for non-owner uploads, collector`md5` keying, WASB-C3****. |

| ****ADR-012**** | GPU-native; TPU deferred; ****Cloud Run+GPU scale-to-zero = serving model****; ****NVDEC**** named as the decode lever | The serving substrate + the 7-step seed scope this expands. |

| ****ADR-013**** | Drop D0; ****GCP sole compute stack****; \$200 D0 credits ? non-compute only | The provider for cloud-serving; D0 Spaces remains only a possible S3 ***storage*** target. |

| ****ADR-014 (new)**** | ****Compute-decoupled serving:**** one pure core, profile-selected compute+storage; stitch-where-metered | This doc. ****Refines**** ADR-011/012, does ****not**** supersede them. |

| PLATFORM_ARCHITECTURE.md | The `ComputeBackend`/`StorageBackend` registries | The abstraction this realizes in code. |

| VIDEO_LOCALITY_MODEL.md / DESKTOP_APP_PLAN.md | L0 fully-local tier / `LocalDesktop` backend | The truly-local profile = these, productionized. |

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

194. assistant (2026-07-08T05:42:31.781Z)

This is important context ? ****D16(b)**** explicitly deferred NVENC encode "behind a config knob, revisit only if M7 cost data justifies relaxing parity." The 2026-07-08 CUJ is exactly that justification, and my design keeps libx264 the default with NVENC opt-in via config ? so it aligns. Let me write the phase-placement doc capturing this:

195. assistant (2026-07-08T05:43:26.970Z)

Phase ? machine placement + compile-on-L4 (#521)

```
> **Status:** ? **SHIPPED (code)** . **Owner:** Avi . **Created:** 2026-07-08 . **Issue:**  
[#521](https://github.com/Khelsutra/badminton-highlight-indexer/issues/521)  
> **Relation:** extends [COMPUTE_DECOUPLED_SERVING](./README.md) (the `ComputeBackend` dispatch  
seam) and **refines D16(b)** (NVENC-encode-deferred-behind-a-config-knob) now that the  
2026-07-08 live CUJ supplies the cost data D16(b) said to revisit on.  
> **Evidence:** `runlogs/DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md` §9  
(2026-07-08 live CUJ, rev `rally-control-plane-00019`).
```

0. TL;DR

The compile/stitch phase used to run **in-process** on the CPU-only 2-vCPU control-plane, where the 2026-07-08 live CUJ measured **~12m55s** for an 11-clip / ~80 s reel from a 1.83 GiB 4K source (~45?90 s per ~7 s clip ? the per-clip branding watermark forces a full 4K re-encode with **no NVENC**). Validation (183 s) runs there too; both serialize on `_LOCAL_COMPUTE_LANE`, so one run pins the control-plane for ~16 min ? the API can't scale past a handful of concurrent users.

#521 makes **where** each pipeline phase runs a config edit, not a code change, and moves the stitch to the **L4 worker (NVENC + 8 vCPU)** the way detection already dispatches:

1. **deployment.phase_placement** ? a per-phase knob (`validation` / `segment` / `compile``) with values `control_plane | cloud_run_l4`. Unset ? **today's** placement, byte-identical.
2. **compile ? L4** ? a `kind='compile'` Cloud Run dispatch mirroring the detect path: the control-plane resolves the reel's segments, stages a compile-spec to the bucket, dispatches `python -m backend.worker compile` on the L4, bucket-polls a done-marker, and returns the same `{status, filepath, download_url}` shape the SPA already expects.
3. **NVENC on the L4** ? `deployment.cloud_stitch_encoder` (default `h264_nvenc`) is forwarded onto the worker's `stitching.encoder`; libx264 stays the everywhere-safe fallback (auto-probed).

The cloud-serving preset opts compile onto the L4 (`phase_placement.compile=cloud_run_l4`).

1. The placement model

`resolve_phase_placement(cfg, phase)` (in `backend/config/placement.py`) is the **single source of truth** ? orchestration and the worker never re-derive the mapping. Precedence: an explicit `deployment.phase_placement.<phase>` wins; otherwise the **legacy default** reproduces pre-#521 behaviour exactly:

phase knob unset (legacy default) meaning
`validation` `control_plane` the authoritative gate (#512/#513) stays on the control-plane
`segment` (detect) `cloud_run_l4` iff <code>compute_target==cloud_run</code> , else <code>control_plane</code> today's DETECT dispatch
`compile` (stitch) `control_plane` the in-process ffmpeg stitch

So a config with **no** `phase_placement` block behaves exactly as before this existed. To relocate a phase, set its knob:

```
```jsonc  
// Move the stitch to the L4 (what the cloud-serving preset does):
```

```

{ "deployment": { "compute_target": "cloud_run",
 "phase_placement": { "compile": "cloud_run_l4" } } }

// Force detection back onto the control-plane on a cloud box (e.g. a debugging run):
{ "deployment": { "compute_target": "cloud_run",
 "phase_placement": { "segment": "control_plane" } } }

// Move validation off the control-plane too (must pair with a cloud segment ? see §4):
{ "deployment": { "compute_target": "cloud_run",
 "phase_placement": { "validation": "cloud_run_l4" } } }
...

`control_plane` runs the phase in-process on the (CPU-only) control-plane service;
`cloud_run_l4`
dispatches it to the GPU/NVENC Cloud Run worker Job.

```

## 2. compile ? L4 dispatch (the flow)

```

`orchestration._execute_compile` consults `resolve_phase_placement(cfg, "compile")`. When it
resolves
to `cloud_run_l4` **and** cloud dispatch is available (`_cloud_available`), it calls
`_dispatch_compile_remote`; otherwise it stitches in-process via the shared `_stitch_reel` core
(unchanged).

`_dispatch_compile_remote` (mirrors `_maybe_dispatch_remote`):

1. **Resolve** the reel's segments HERE ? the control-plane owns the DB (`_resolve_compile` ?
 `(video, kept, out_name)`), applying the `stitching.min_shot_count` floor as before.
2. **Stage a compile-spec** to `/_compile/<token>/spec.json` ? the resolved
 `time_pairs`, the control-plane's resolved `stitching` config (watermark + padding, with
 `encoder`
 swapped to `cloud_stitch_encoder`), the source `video_uri`, `video_id`, `output_name`,
 `locale`.
 The worker therefore needs **no control-plane DB** ? everything rides the spec.
3. **Dispatch** a `compile` Cloud Run Job execution (`CompileJobSpec` ?
 `CloudRunCompute.run_compile`)
 running `python -m backend.worker compile --spec-uri ? --done-uri ? --set
 storage.output_root=?`,
 forwarding the bucket coordinates (the worker Job has no `DEPLOYMENT_PROFILE`).
4. **Bucket-poll** the done-marker at `/_compile/<token>.done` ? the SAME
 `_await_marker` machinery `run_infer` uses (crash-fast on a terminal-without-marker
 execution,
 rescue-or-cancel on a poll timeout). Distinct `_compile/` prefix so a compile and a detection
 for
 one bucket never collide on a marker.
5. The worker (`cmd_compile`) reads the spec, stitches the reel via the **shared `_stitch_reel`
 core** (no forked pipeline), mirror-uploads it to
 `/highlights/<video_id>/<name>`,
 and writes the marker (rc + `download_url` + `reel_uri`).
6. The control-plane returns `{status, filepath, download_url, video_id}` ? the reel is already
 in the
 bucket, so `/api/highlights` serves it via the existing signed-URL 307. On a worker rc?0 ?
 `status='failed'` with the worker's message (clean SPA verdict, never a crash).

Fallback: if compile is placed on the L4 but the box isn't actually cloud-wired (no output
bucket
/ Cloud Run env), `_dispatch_compile_remote` logs and stitches in-process ? a slow reel beats no
reel.

```

## 3. NVENC on the L4 ? and the D5/D16(b) parity note



The dominant term in the CUJ stitch was the **per-clip 4K re-encode** the branding watermark forces (~67 s/clip on 2 vCPU with libx264). The L4 has a hardware H.264/HEVC encoder, so #521 adds an encoder knob:

```
- `StitchingConfig.encoder` ? `libx264` (default, CPU, works everywhere incl. CI) | `h264_nvenc`
|
 `hevc_nvenc`. `VideoCompiler` probes the requested encoder once per run (a 1-frame lavfi
test
 encode) and falls back to libx264 if this ffmpeg build can't run it, so a misconfig / non-
GPU
 box never breaks. All clips share the one chosen encoder ? the concat stays -c copy`.
- `deployment.cloud_stitch_encoder` (default h264_nvenc`) ? what the compile-dispatch forwards
onto
 the worker's stitching.encoder` for the L4 run. Set it to libx264` to run the L4 stitch on
CPU.
- The control-plane's own local stitch always uses stitching.encoder` (libx264 by default)
? the
 CP has no GPU.
```

**Parity note (honest):** [README.md](./README.md) **D5/D16(b)** kept stitch **encode = libx264** to protect a byte-exact stitch parity guarantee, and explicitly **deferred NVENC** encode behind a config knob, "revisit only if M7 cost data justifies relaxing parity." The 2026-07-08 CUJ (~13 min stitch, a hard scaling wall past ~5 concurrent users) **is** that cost data. #521 lands exactly the deferred shape ? NVENC **behind a config knob**, libx264 the default, config-reversible per phase. Enabling NVENC relaxes byte-identity **at the reel ENCODE layer only** (a cosmetic, shareable artifact); the **WINDOW/rally-level parity** (the analytical output) is untouched ? detection and windowing are unchanged. If bit-exact reels are ever required, set `cloud_stitch_encoder=libx264``.

> Follow-up (tracked, not in #521's first cut): a **pre-rendered / overlay-filter watermark** so the  
> stitch isn't a full re-encode at all (the deeper win beyond swapping the encoder).

---

#### 4. Validation placement + the "nobody validates" guard

`phase_placement.validation`` moves the CPU-bound validator suite off the 2-vCPU control-plane too, via the existing `--skip-validation`` lever (#513):

```
- control_plane` (default): the control-plane runs the validators and is the authoritative
gate;
 the worker skips its redundant re-validation (skip_validation=True`) ? today's #512/#513
 behaviour, byte-identical.
- cloud_run_l4`: when detection actually dispatches to the L4, the control-plane
delegates
 ? it does NOT run the validator suite (_execute_processing` skips it), and the dispatch
 carries
 no --skip-validation`, so the worker validates and rejects a bad clip via a non-zero
 rc
 (surfaced as a failed job).
```

**Footgun guard:** a clip can only be validated on the worker if the worker also **detects** it there.

Two protections:

```

1. **Config-load guard** (`DeploymentConfig` model-validator): `validation=cloud_run_l4`
requires the
 effective segment placement to also be `cloud_run_l4` (explicit, or via
`compute_target=cloud_run`)
 ? else config load fails with an actionable error.
2. **Runtime guard** (`_will_detect_on_cloud`): if a per-request `compute='local'` forces
detection
 in-process (no worker), the control-plane **still validates** ? never leave a clip
unvalidated
 because nobody ran the gate.

```

---

## 5. Files & tests

Area	Code	Tests
Placement resolver + config	`backend/config/placement.py`, `backend/config/models.py` (`PhasePlacementConfig`, `DeploymentConfig.phase_placement` / `cloud_stitch_encoder`, guard)	`tests/test_phase_placement.py`, `tests/test_config_deployment.py`
Preset	`backend/config/presets.py` (`cloud-serving` ? `phase_placement.compile=cloud_run_l4`)	`tests/test_phase_placement.py`
compile dispatch	`backend/orchestration.py` (`_execute_compile`, `_dispatch_compile_remote`, `_stitch_reel`, `_will_detect_on_cloud`)	`tests/test_compile_dispatch.py`
compute backend	`backend/compute/base.py` (`CompileJobSpec`, `run_compile`), `backend/compute/cloud_run.py` (`run_compile`, `_await_marker`)	`tests/test_compute_backend.py`
worker	`backend/worker/__main__.py` (`compile` subcommand, `cmd_compile`)	`tests/test_worker.py`
NVENC encoder	`backend/pipeline/stitcher/compiler.py` (`encoder`, `_select_encoder`, `_probe_encoder_available`)	`tests/test_compiler.py`
validation delegation	`backend/orchestration.py` (`_maybe_dispatch_remote` skip flag, `_execute_processing`)	`tests/test_api_cloud_dispatch.py`

All GPU-free + cloud-free: a fake `CloudRunJobClient` simulates the worker (reads the staged spec, writes the reel + done-marker into a local "bucket"), mirroring  
`tests/test\_api\_cloud\_dispatch.py`.

### 196. user (2026-07-08T05:43:27.007Z)

File created successfully at: C:\kg-wt521\docs\COMPUTE\_DECOUPLED\_SERVING\PHASE\_PLACEMENT.md  
(file state is current in your context ? no need to Read it back)

### 197. assistant (2026-07-08T05:43:29.786Z)

Now add a pointer to the new doc from the README, right after the honesty note:

### 198. user (2026-07-08T05:43:34.375Z)

<tool\_use\_error>File has not been read yet. Read it first before writing to it.</tool\_use\_error>

### 199. user (2026-07-08T05:43:38.900Z)

```

9
10 ---
11
12 > **? Northstar (D14):** a **mobile, app-like** flow ? point at a **Google Drive** video
? **fully cloud-native** run ? the stitched, **ready-to-share** reel lands **back in Drive, next
to the source**. Two design promises around it: a user can **switch local?cloud anytime** (D13,
a headline feature), and execution is **never a surprise** ? auto-detect **suggests**, the user
confirms (D15).

```

### 200. user (2026-07-08T05:43:41.708Z)

```
8 > **Honesty note:** claims tagged `[verified]` (checked against code by the design
workflow `wf_56274ff0`), `[design]` (proposed), `[researched]`.
9
10 ---
```

#### 201. user (2026-07-08T05:43:46.740Z)

The file C:\kg-wt521\docs\COMPUTE\_DECOUPLED\_SERVING\README.md has been updated successfully.  
(file state is current in your context ? no need to Read it back)

#### 202. assistant (2026-07-08T05:43:50.396Z)

Now re-run the doc gates and the full set of touched test files to confirm everything's green:

#### 203. user (2026-07-08T05:43:53.429Z)

```
=== md_links ===
[check-md-links] clean ? no broken internal links in live docs.
=== leak_scan ===
[leak-scan] clean ? no MD5 / personal-path / private-name leaks in the code surface.
=== ruff ===
All checks passed!
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

#### 204. user (2026-07-08T05:44:44.227Z)

```
tests/test_api_cloud_dispatch.py::test_full_api_process_job_result_carries_cloud_provenance
C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\site-
packages\google\auth_default.py:113: UserWarning: Your application has authenticated using end
user credentials from Google Cloud SDK without a quota project. You might receive a "quota
exceeded" or "API not enabled" error. See the following page for troubleshooting:
https://cloud.google.com/docs/authentication/adc-troubleshooting/user-creds.
warnings.warn(_CLOUD_SDK_CREDENTIALS_WARNING)

-- Docs: https://docs.pytest.org/en/stable/how-to/capture-warnings.html
225 passed, 8 warnings in 44.19s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

#### 205. assistant (2026-07-08T05:44:57.688Z)

All 225 pass. Let me run the complete suite once more with coverage to confirm the whole thing  
is green after the budget fix (honoring the pre-LGTM full-suite gate):

#### 206. user (2026-07-08T05:48:36.692Z)

```
training\gen0\harness.py 130 9 93%

TOTAL 12663 1810 86%
Required test coverage of 70% reached. Total coverage: 85.71%
1875 passed, 8 skipped, 25 warnings in 213.65s (0:03:33)
=====
RESULT: PASS - all tests green
=====
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

#### 207. assistant (2026-07-08T05:48:49.355Z)

Full suite green (1875 passed, 85.71% coverage). Let me re-confirm mypy and check the worktree's  
diff before committing (branching-safety: verify only my commits, stage explicit paths):

#### 208. user (2026-07-08T05:48:56.168Z)

```
=== mypy ===
```

```
backend\providers__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 150 source files
=== branch ===
feat/521-compile-on-l4
=== ahead of origin/master (should be empty) ===
=== changed files ===
M backend/compute/__init__.py
M backend/compute/base.py
M backend/compute/cloud_run.py
M backend/config/__init__.py
M backend/config/models.py
M backend/config/presets.py
M backend/orchestration.py
M backend/pipeline/stitcher/compiler.py
M backend/worker/__main__.py
M docs/COMPUTE_DECOUPLED_SERVING/README.md
M tests/test_api_cloud_dispatch.py
M tests/test_compiler.py
M tests/test_compute_backend.py
M tests/test_config_deployment.py
M tests/test_worker.py
?? backend/config/placement.py
?? docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md
?? tests/test_compile_dispatch.py
?? tests/test_phase_placement.py
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

#### 209. assistant (2026-07-08T05:49:07.247Z)

mypy clean, branch clean (no sibling commits), `main.py` correctly nets to zero. Let me stage explicit paths (never `-A`) and commit:

#### 210. user (2026-07-08T05:49:10.960Z)

```
warning: in the working copy of 'backend/config/placement.py', LF will be replaced by CRLF the
next time Git touches it
warning: in the working copy of 'docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md', LF will be
replaced by CRLF the next time Git touches it
warning: in the working copy of 'tests/test_compile_dispatch.py', LF will be replaced by CRLF
the next time Git touches it
warning: in the working copy of 'tests/test_phase_placement.py', LF will be replaced by CRLF the
next time Git touches it
=== staged ===
M backend/compute/__init__.py
M backend/compute/base.py
M backend/compute/cloud_run.py
M backend/config/__init__.py
M backend/config/models.py
A backend/config/placement.py
M backend/config/presets.py
M backend/orchestration.py
M backend/pipeline/stitcher/compiler.py
M backend/worker/__main__.py
A docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md
M docs/COMPUTE_DECOUPLED_SERVING/README.md
M tests/test_api_cloud_dispatch.py
A tests/test_compile_dispatch.py
M tests/test_compiler.py
M tests/test_compute_backend.py
M tests/test_config_deployment.py
A tests/test_phase_placement.py
M tests/test_worker.py
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 211. assistant (2026-07-08T05:49:16.143Z)

All 19 paths staged (the CRLF warnings are harmless line-ending normalization). Now commit:

### 212. user (2026-07-08T05:49:30.634Z)

```
[feat/521-compile-on-l4 8866ffe] perf(serving): run compile/stitch on the L4 worker; config-
driven phase?machine placement (#521)
 19 files changed, 1710 insertions(+), 55 deletions(-)
 create mode 100644 backend/config/placement.py
 create mode 100644 docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md
 create mode 100644 tests/test_compile_dispatch.py
 create mode 100644 tests/test_phase_placement.py
=== verify ===
8866ffe perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521)
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 213. assistant (2026-07-08T05:49:34.632Z)

Clean single commit. Now push and open the PR against master:

### 214. user (2026-07-08T05:49:38.643Z)

```
remote:
remote: Create a pull request for 'feat/521-compile-on-l4' on GitHub by visiting:
remote: https://github.com/Khelsutra/badminton-highlight-indexer/pull/new/feat/521-compile-
on-l4
remote:
branch 'feat/521-compile-on-l4' set up to track 'origin/feat/521-compile-on-l4'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
 * [new branch] feat/521-compile-on-l4 -> feat/521-compile-on-l4
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 215. assistant (2026-07-08T05:49:45.626Z)

Let me check the PR template, then open the PR:

### 216. user (2026-07-08T05:49:46.730Z)

```
<!-- Describe the change and link the issue it closes. -->
```

## Summary

## Checks

- [ ] `python run\_all\_tests.py` green; `ruff check .` + `mypy` clean; coverage not lowered.
- [ ] Internal markdown links valid (`python -m tools.check\_md\_links`).

## Launch / "bit-flip" PRs only

A PR that flips a default / "launch bit" so a tracked project reaches terminal state MUST also run

the `docs/DOC\_STATUS.md` \$2 archival due-diligence on that project's docs, **in this same PR**:

- [ ] Verified each doc's claims against the code (tagged load-bearing facts `[verified]`).
- [ ] Flipped the doc Status ? DELIVERED/DONE/REJECTED + a completion note (what shipped, PR, paths).
- [ ] Updated all inbound references (`docs/README.md`, `CLAUDE.md`, `docs/NEXT\_STEPS.md`, cross-doc links).
- [ ] Kept the ADR forward/back lineage.
- [ ] `git mv` the terminal docs to `docs/archives/?` and recorded the move in `DOC\_STATUS.md` \$3.
- [ ] Fixed any internal-link breakage from the move (`link-check` CI job green).

\_(Not a launch/bit-flip PR? Delete this section.)\_

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

#### 217. user (2026-07-08T05:50:11.329Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/525>

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

#### 218. assistant (2026-07-08T05:50:26.128Z)

PR #525 created. Let me verify it targets master and CI is running:

#### 219. user (2026-07-08T05:50:31.007Z)

```
{"base":"master","head":"feat/521-compile-on-l4","number":525,"state":"OPEN","title":"perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine placement (#521)","url":"https://github.com/Khelsutra/badminton-highlight-indexer/pull/525"}
```

=== checks (may be pending) ===

Golden fixtures cross-OS (ubuntu-latest) pending 0

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160128>

Ruff lint pass 11s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160091>

Golden fixtures cross-OS (macos-latest) pending 0 <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160157>

Leak scan (PII / attribution guard) pass 8s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160101>

Markdown link check (live docs) pass 6s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160086>

Built wheel smoke (install + launch) pending 0 <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160154>

Dependency license scan (no copyleft in runtime deps) pending 0

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160080>

Full test suite (offline, mocked) pending 0 <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160106>

Golden fixtures cross-OS (windows-latest) pending 0

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160184>

Mypy (lenient baseline) pending 0 <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160107>

Syntax sweep + import smoke (no GPU stack) pending 0

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160088>

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

#### 220. assistant (2026-07-08T05:50:41.795Z)

PR #525 is open against master; early gates (ruff, leak-scan, md-links) already pass. Let me watch CI to completion:

#### 221. user (2026-07-08T05:53:19.523Z)

Full test suite (offline, mocked) pass 2m50s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160106>

Golden fixtures cross-OS (macos-latest) pass 25s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160157>

Golden fixtures cross-OS (ubuntu-latest) pass 27s

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160128>

Golden fixtures cross-OS (windows-latest) pass 53s

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/28920721803/job/85797160184>

```

Leak scan (PII / attribution guard) pass 8s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28920721803/job/85797160101
Markdown link check (live docs) pass 6s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28920721803/job/85797160086
Mypy (lenient baseline) pass 34s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28920721803/job/85797160107
Ruff lint pass 11s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28920721803/job/85797160091
Syntax sweep + import smoke (no GPU stack) pass 20s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28920721803/job/85797160088
Built wheel smoke (install + launch) pass 1m4s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28920721803/job/85797160154
Dependency license scan (no copyleft in runtime deps) pass 34s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28920721803/job/85797160080
Full test suite (offline, mocked) pass 2m50s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28920721803/job/85797160106
Golden fixtures cross-OS (macos-latest) pass 25s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28920721803/job/85797160157
Golden fixtures cross-OS (ubuntu-latest) pass 27s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28920721803/job/85797160128
Golden fixtures cross-OS (windows-latest) pass 53s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28920721803/job/85797160184
Leak scan (PII / attribution guard) pass 8s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28920721803/job/85797160101
Markdown link check (live docs) pass 6s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28920721803/job/85797160086
Mypy (lenient baseline) pass 34s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28920721803/job/85797160107
Ruff lint pass 11s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28920721803/job/85797160091
Syntax sweep + import smoke (no GPU stack) pass 20s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28920721803/job/85797160088
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

## 222. assistant (2026-07-08T05:53:27.609Z)

All 11 CI checks green. Let me file the deferred Part 3 (cheaper watermark) as a tracked follow-up and record a brief session-handoff memory:

## 223. user (2026-07-08T05:53:42.799Z)

Noted (position 1, task\_id: task\_217edf9c). A chip is showing for the user ? they can start it in a fresh worktree with one click, or dismiss it. If this suggestion becomes stale or superseded, call dismiss\_task with this task\_id. Currently pending: task\_217edf9c "Cheaper reel watermark (avoid full re-encode)". Continue your current work.

## 224. assistant (2026-07-08T05:53:58.391Z)

---

```

name: pr-525-compile-on-l4
description: PR #525 (#521) shipped config-driven phase?machine placement + compile-on-L4 +
NVENC; Part 3 (cheaper watermark) deferred
metadata:
 type: project

```

**\*\*PR [#525](https://github.com/Khelsutra/badminton-highlight-indexer/pull/525)\*\*** (closes #521) ? opened 2026-07-08, all 11 CI checks green, awaiting owner review/merge.

Drove the compile/stitch off the CPU-only 2-vCPU control-plane (2026-07-08 live CUJ: ~12m55s for an 11-clip 4K reel ? §9 of the validation-latency runlog) to the L4 worker.

**\*\*What shipped:\*\***

- `deployment.phase\_placement` (`validation`/`segment`/`compile` ? `control\_plane|cloud\_run\_l4`), resolved by `backend/config/placement.py::resolve\_phase\_placement`. Unset ? legacy placement, byte-identical. Config-load guard blocks validation-on-L4 without a cloud segment.
- `kind='compile'` Cloud Run dispatch (`orchestration.\_dispatch\_compile\_remote`) mirroring `\_maybe\_dispatch\_remote`; worker `compile` subcommand (`backend/worker/\_\_main\_\_.py::cmd\_compile`) stitches via the shared `orchestration.\_stitch\_reel` core (no forked pipeline, no worker DB ? a staged bucket spec carries resolved time-pairs + stitching config).
- NVENC: `StitchingConfig.encoder` + `deployment.cloud\_stitch\_encoder` (default `h264\_nvenc`); `VideoCompiler.\_select\_encoder` probes + falls back to libx264. Lands the D16(b)-deferred "NVENC behind a config knob".
- cloud-serving preset sets `phase\_placement.compile=cloud\_run\_l4` (validation stays on control\_plane by default even in cloud-serving).

**\*\*Deferred (tracked):\*\*** Part 3 = cheaper watermark so the stitch isn't a full re-encode (pre-rendered overlay, or watermark-at-concat so per-clip trims stream-copy). Flagged in `docs/COMPUTE\_DECOUPLED\_SERVING/PHASE\_PLACEMENT.md` §3.

**\*\*Not verified (owner-side):\*\*** real L4 NVENC stitch wall ? deploy was out of scope; no `gcloud run deploy` made. Worktree at `C:/kg-wt521` (branch `feat/521-compile-on-l4`).

#### 225. user (2026-07-08T05:53:58.427Z)

File created successfully at: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\pr-525-compile-on-l4.md (file state is current in your context ? no need to Read it back)

#### 226. user (2026-07-08T05:54:03.832Z)

pr-525-compile-on-l4.md

#### 227. assistant (2026-07-08T05:54:07.961Z)

## Memory index

- [PR #525 ? compile-on-L4 (#521)](pr-525-compile-on-l4.md) ? config-driven phase?machine placement + L4 NVENC stitch; Part 3 (cheaper watermark) deferred.

#### 228. user (2026-07-08T05:54:07.997Z)

File created successfully at: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\MEMORY.md (file state is current in your context ? no need to Read it back)

#### 229. assistant (2026-07-08T05:54:28.996Z)

Done. All three parts of #521 are implemented, tested, and shipped as **\*\*[PR #525](https://github.com/Khelsutra/badminton-highlight-indexer/pull/525)\*\*** (against `master`, all **\*\*11 CI checks green\*\***).

## What landed

**\*\*1. Compile/stitch dispatched to the L4 worker.\*\*** A `kind='compile'` Cloud Run dispatch mirroring `\_maybe\_dispatch\_remote`: the control-plane resolves the reel's segments, stages a compile-spec (resolved time-pairs + stitching config + source URI) to the bucket, dispatches `python -m backend.worker compile` on the L4, bucket-polls a done-marker (same crash-fast/timeout-rescue machinery as detect), and returns the same `{status, filepath, download\_url}` shape the SPA expects. The worker stitches via the **\*\*shared `\_stitch\_reel` core\*\***



(no forked pipeline, no worker DB). Control-plane stitch stays available as a config option.

**\*\*2. Config-driven phase?machine placement.\*\*** ``deployment.phase_placement`` (``validation`/`segment`/`compile` ? `control_plane | cloud_run_l4``), resolved by `[`backend/config/placement.py`](backend/config/placement.py)`. **\*\*Unset ? today's placement, byte-identical.\*\*** A config-load guard rejects the "nobody validates" footgun. The cloud-serving preset sets ``compile` ? cloud_run_l4``. Flipping the knob moves validation/segment/compile between machines with no code change.

**\*\*3. NVENC on the L4.\*\*** ``StitchingConfig.encoder`` + ``deployment.cloud_stitch_encoder`` (default ``h264_nvenc``) forwarded to the worker; ``VideoCompiler`` probes it once and **\*\*falls back to libx264\*\*** on CI/non-GPU boxes. This lands the D16(b)-deferred "NVENC-encode behind a config knob" now that the CUJ supplies the justifying cost data ? libx264 stays the parity-safe default; window/rally-level parity is untouched.

## Verification (pre-LGTM gate, run locally)

- **\*\*Full suite: 1875 passed, 8 skipped\*\***; ``ruff check .`` + ``mypy backend training`` clean; ``tools.leak_scan`` + ``tools.check_md_links`` clean; **\*\*coverage 85.7%\*\*** (floor 70). All GPU-free/cloud-free via a fake ``CloudRunJobClient``.
- New tests: ``test_phase_placement.py``, ``test_compile_dispatch.py``; additions to ``test_compute_backend.py``, ``test_worker.py``, ``test_compiler.py``, ``test_api_cloud_dispatch.py``.

## Honest limits

- **\*\*Deploy was out of scope\*\*** ? no ``gcloud run deploy`/traffic change` was made, so the *\*real L4 NVENC stitch wall\** isn't measured here (that needs an owner-driven deploy + CUJ). The dispatch plumbing is proven end-to-end with mocks.
- **\*\*Part 3 deferred\*\*** as the task allowed: the NVENC change hardware-accelerates the *\*encode\**, but each clip is still fully re-encoded because the watermark filter forces it. The deeper win (pre-rendered overlay / watermark-at-concat so trims stream-copy) is filed as a tracked follow-up chip and noted in ``docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md`` §3.

Work was done in an isolated worktree (``C:/kg-wt521``, branch ``feat/521-compile-on-l4``) ? the shared checkout and sibling branches were never touched.

### 230. user (2026-07-08T05:54:31.372Z)

please open a PR with the changes and ensure thorough testing and use the rally-worker setup to test on real videos if possible

### 231. assistant (2026-07-08T05:55:30.719Z)

The PR is already open (#525) ? I'll confirm it, then focus on the new asks: **\*\*thorough testing\*\*** (including an adversarial review of my own diff) and **\*\*real-video testing via the rally-worker setup\*\***. Let me start with parallel recon: PR state, available real videos in GCS, local ffmpeg/NVENC, the deploy tooling, and the currently-deployed worker image.

### 232. user (2026-07-08T05:55:34.790Z)

=== PR 525 ===

```
{
 "checks": [
 {
 "concl": "SUCCESS",
 "name": "Built wheel smoke (install + launch)"
 },
 {
 "concl": "SUCCESS",
 "name": "Dependency license scan (no copyleft in runtime deps)"
 },
 {
 "concl": "SUCCESS",
 "name": "Full test suite (offline, mocked)"
 },
 {
 "concl": "SUCCESS",
 "name": "Golden fixtures cross-OS (macos-latest)"
 },
 {
 "concl": "SUCCESS",
 "name": "Golden fixtures cross-OS (ubuntu-latest)"
 },
 {
 "concl": "SUCCESS",
 "name": "Golden fixtures cross-OS (windows-latest)"
 },
 {
 "concl": "SUCCESS",
 "name": "Leak scan (PII / attribution guard)"
 },
 {
 "concl": "SUCCESS",
 "name": "Markdown link check (live docs)"
 },
 {
 "concl": "SUCCESS",
 "name": "Mypy (lenient baseline)"
 },
 {
 "concl": "SUCCESS",
 "name": "Ruff lint"
 },
 {
 "concl": "SUCCESS",
 "name": "Syntax sweep + import smoke (no GPU stack)"
 }
],
 "mergeStateStatus": "CLEAN",
 "mergeable": "MERGEABLE",
 "state": "OPEN"
}
```

Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

### 233. user (2026-07-08T05:55:41.168Z)

```
=== gcloud auth/project ===
[accessibility]
screen_reader = False
[billing]
quota_project = gen-lang-client-0306026826
[core]
account = avi.dullu@gmail.com
disable_usage_reporting = False
project = gen-lang-client-0306026826

Your active configuration is: [default]
=== corpus videos (small ones) ===
gs://khelsutra-rally-corpus/videos/GX010128.MP4
gs://khelsutra-rally-corpus/videos/GX010142_1080p.mp4
gs://khelsutra-rally-corpus/videos/GX010143_1080p.mp4
gs://khelsutra-rally-corpus/videos/GX010147_1080p.mp4
gs://khelsutra-rally-corpus/videos/GX030094_proxy.mp4
gs://khelsutra-rally-corpus/videos/mahadevpura-1.MP4
gs://khelsutra-rally-corpus/videos/mahadevpura_1.mp4
gs://khelsutra-rally-corpus/videos/mahadevpura_smoke_180s.mp4
gs://khelsutra-rally-corpus/videos/SHANAN/
gs://khelsutra-rally-corpus/videos/golden/
gs://khelsutra-rally-corpus/videos/golden_1080p/
```

#### **234. user (2026-07-08T05:55:42.925Z)**

```
=== ffmpeg version ===
ffmpeg version 8.1.1-full_build-www.gyan.dev Copyright (c) 2000-2026 the FFmpeg developers
built with gcc 15.2.0 (Rev13, Built by MSYS2 project)
=== nvenc encoders present? ===
V....D av1_nvenc NVIDIA NVENC av1 encoder (codec av1)
V....D libx264 libx264 H.264 / AVC / MPEG-4 AVC / MPEG-4 part 10 (codec h264)
V....D libx264rgb libx264 H.264 / AVC / MPEG-4 AVC / MPEG-4 part 10 RGB (codec h264)
V....D h264_amf AMD AMF H.264 Encoder (codec h264)
V....D h264_d3d12va D3D12VA h264 encoder (codec h264)
V....D h264_mf H264 via MediaFoundation (codec h264)
```

#### **235. user (2026-07-08T05:55:43.686Z)**

```
=== deploy tooling ===
deploy/cloudrun:
cloudbuild.yaml
Dockerfile
gen_service_yaml.py
post_deploy_smoke.ps1
README.md

deploy/gcp:
buckets.env
create_gpu_vm.sh
nightly_regression
ops-agent-startup.sh
provision.sh
README.md
teardown.sh
=== build/deploy scripts ===
ls: cannot access 'deploy/cloudrun/*.sh': No such file or directory
deploy/gcp/create_gpu_vm.sh
deploy/gcp/ops-agent-startup.sh
deploy/gcp/provision.sh
deploy/gcp/teardown.sh
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

#### **236. assistant (2026-07-08T05:55:56.142Z)**

PR #525 is confirmed **\*\*OPEN, MERGEABLE, all 11 checks green\*\***. Local ffmpeg has NVENC-family encoders but **\*\*no `h264\_nvenc`\*\*** (this is an AMD box) ? so I can test the libx264 fallback path against a real probe. Let me check the deployed worker image, video sizes, and the build flow:

### 237. user (2026-07-08T05:56:04.183Z)

```
=== rally-worker Job image ===
asia-southeast1-docker.pkg.dev/gen-lang-client-0306026826/rally/rally-worker:latest
=== candidate video sizes ===
 767766440 2026-06-22T13:42:54Z gs://khehsutra-rally-corpus/videos/GX010142_1080p.mp4
TOTAL: 1 objects, 767766440 bytes (732.20MiB)
 90194321 2026-06-20T14:20:32Z gs://khehsutra-rally-corpus/videos/mahadevpura_smoke_180s.mp4
TOTAL: 1 objects, 90194321 bytes (86.02MiB)
1429944539 2026-06-20T14:27:35Z gs://khehsutra-rally-corpus/videos/GX030094_proxy.mp4
TOTAL: 1 objects, 1429944539 bytes (1.33GiB)
```

### 238. user (2026-07-08T05:56:04.819Z)

```
=== cloudbuild.yaml ===
```

**Cloud Build config for the Cloud Run GPU worker image (COMPUTE\_DECOUPLED\_**

**#**  
**`gcloud builds submit --tag ? -f ?` does NOT work ? `gcloud builds submit` has no `-f`**  
**non-root Dockerfile needs a build CONFIG like this one. Build with:**

**#**  
**gcloud builds submit --config=deploy/cloudrun/cloudbuild.yaml --project <PROJECT**

**#**  
**Override the region/repo/image with --substitutions=\_REGION=?,\_REPO=?,\_IMAGE=?**  
**PROJECT\_ID is provided by Cloud Build automatically.**

```
substitutions:
 _REGION: asia-southeast1 # L4 GPU is region-limited (asia-southeast1 / europe-west1 /
europe-west4 / us-central1 / us-east4)
 _REPO: rally
 _IMAGE: rally-worker
steps:
- name: 'gcr.io/cloud-builders/docker'
 args:
 - build
 - -t
 - '${_REGION}-docker.pkg.dev/${PROJECT_ID}/${_REPO}/${_IMAGE}:latest'
 - -f
 - deploy/cloudrun/Dockerfile
 - .
images:
- '${_REGION}-docker.pkg.dev/${PROJECT_ID}/${_REPO}/${_IMAGE}:latest'
options:
 machineType: E2_HIGHCPU_8 # the CUDA + micromamba image is heavy; a bigger builder keeps it
~8?10 min
timeout: 3600s
Shell cwd was reset to C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-indexer
```

### 239. assistant (2026-07-08T05:56:40.721Z)

PR is fully green and mergeable. Now the important findings for the L4 test: the deployed `rally-worker` runs `:latest`, built from an **\*\*old commit without my `compile` subcommand\*\*** ? so a real L4 compile test requires **\*\*building a new image\*\***, and `:latest` is the **\*shared live tag\*** (control-plane + worker both use it). That's a traffic-affecting deploy I must not do unilaterally.

Let me do the strongest tests I **\*can\*** do safely first: a **\*\*real-video end-to-end worker compile\*\*** (real ffmpeg, real clips) and a **\*\*real-GCS dispatch round-trip\*\***. Let me pull a real

clip:

#### 240. user (2026-07-08T05:56:55.910Z)

pulling real 180s clip...

.....

Average throughput: 21.7MiB/s

=== source probe ===

codec\_name=h264

width=1920

height=1080

r\_frame\_rate=30000/1001

codec\_name=aac

r\_frame\_rate=0/0

codec\_name=unknown

r\_frame\_rate=0/0

duration=180.046167

size=90194321

#### 241. assistant (2026-07-08T05:57:13.529Z)

Real 1080p/180s H.264 clip. Now let me run the **worker compile** subcommand end-to-end on the real video ? with `encoder=h264_nvenc` in the spec (as the L4 dispatch forwards), which will exercise the real probe ? libx264 fallback (no NVIDIA GPU here) plus a real watermark burn:

#### 242. user (2026-07-08T05:57:20.482Z)

video\_id=05e0e577029daafc425ca70c0c4f0a3

=== spec written, video\_id in reel path ===

#### 243. user (2026-07-08T05:57:42.128Z)

=== RUN worker compile on the REAL 180s clip ===

11:27:34 [INFO] backend.pipeline.stitcher.compiler: Source video duration: 180.05s

11:27:34 [INFO] backend.pipeline.stitcher.compiler: Trimming 4 segments (begin\_padding=0.5s, end\_padding=0.5s)...

11:27:34 [INFO] backend.pipeline.stitcher.compiler: [watermark] applying branding overlay to each clip.

11:27:35 [WARNING] backend.pipeline.stitcher.compiler: [compile] encoder 'h264\_nvenc' is not runnable in this ffmpeg build ? falling back to libx264 (CPU). On the L4 worker this should be available.

11:27:35 [INFO] backend.pipeline.stitcher.compiler: Trimming Segment #1/4 [5.00s - 12.00s] -> [4.500s - 12.500s] (duration: 8.00s)

11:27:36 [INFO] backend.pipeline.stitcher.compiler: Trimming Segment #2/4 [40.00s - 48.00s] -> [39.500s - 48.500s] (duration: 9.00s)

11:27:38 [INFO] backend.pipeline.stitcher.compiler: Trimming Segment #3/4 [90.00s - 98.00s] -> [89.500s - 98.500s] (duration: 9.00s)

11:27:39 [INFO] backend.pipeline.stitcher.compiler: Trimming Segment #4/4 [150.00s - 158.00s] -> [149.500s - 158.500s] (duration: 9.00s)

11:27:41 [INFO] backend.pipeline.stitcher.compiler: Successfully trimmed 4/4 clips. Preparing to stitch...

11:27:41 [INFO] backend.pipeline.stitcher.compiler: Stitching clips together...

11:27:41 [INFO] backend.pipeline.stitcher.compiler: Dynamic highlight reel saved:

C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer/75d3d121-f240-4293-afa5-b436a379ad31/scratchpad/compiletest/wscratch\smoke-reel.mp4 (21.8 MB)

11:27:41 [INFO] backend.worker: [worker] wrote compile marker ->

C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-

indexer/75d3d121-f240-4293-afa5-b436a379ad31/scratchpad/compiletest/bucket/\_compile/tok.done

[worker] compile: video\_id=05e0e577029daafc425ca70c0c4f0a3 clips=4 encoder=h264\_nvenc

out=smoke-reel.mp4

[worker] compile: stitching

[worker] compile: finished

real 0m10.161s  
user 0m0.123s  
sys 0m0.123s  
RC=0

Shell cwd was reset to C:\Users\avidu\Projects\khelesutra-guru\badminton-highlight-indexer

#### 244. assistant (2026-07-08T05:57:49.224Z)

The worker compile ran end-to-end on the real clip in ~10s: probed `h264\_nvenc` ? \*\*fell back to libx264\*\* (fallback works on real ffmpeg), trimmed 4/4 real clips with watermark, stitched a 21.8 MB reel, wrote the marker, RC=0. Let me verify the mirrored reel + marker + that the reel is a valid playable mp4:

#### 245. user (2026-07-08T05:57:54.868Z)

```
=== done-marker ===
{"video_id": "05e0e577029daafc425ca70c0c4f0a3", "returncode": 0, "status": "success",
"download_url": "/api/highlights/05e0e577029daafc425ca70c0c4f0a3/smoke-reel.mp4", "reel_uri":
"", "message": ""}
=== mirrored reel in bucket (highlights/<video_id>/) ===
ls: cannot access 'C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelesutra-
guru-badminton-highlight-indexer/75d3d121-f240-4293-afa5-
b436a379ad31/scratchpad/compiletest/bucket/highlights/05e0e577029daafc425ca70c0c4f0a3/': No
such file or directory
=== reel is a valid mp4? (ffprobe the mirrored copy) ===
C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelesutra-guru-badminton-
highlight-indexer/75d3d121-f240-4293-afa5-
b436a379ad31/scratchpad/compiletest/bucket/highlights/05e0e577029daafc425ca70c0c4f0a3/smoke-
reel.mp4: No such file or directory
=== expected reel duration ~ sum of padded clips (8+9+9+9 = 35s) ===
```

#### 246. assistant (2026-07-08T05:58:32.987Z)

The reel was produced (21.8 MB) but the mirror was correctly skipped ? my local-path `output\_root` has no `://`, so `reel\_object\_uri` returns `""` (that's the intended "truly-local ? serve the local file" path). On the real L4 the `output\_root` is `gs://`, so the mirror fires. Let me verify the produced reel is valid, then run the \*\*real-GCS round-trip\*\* (fetch source from GCS + upload reel to GCS + marker to GCS) to exercise the exact storage plumbing the L4 uses:

#### 247. user (2026-07-08T05:58:49.018Z)

```
=== the produced reel (local output_dir) is a valid mp4 ===
codec_name=h264
width=1920
height=1080
codec_name=aac
codec_name=unknown
duration=35.089000
size=22863973
=== gcs lib present locally? ===
google-cloud-storage OK
=== write access to serving bucket scratch prefix? ===
```

```
.....
probe 1783490318
```

#### 248. assistant (2026-07-08T05:59:00.525Z)

Reel is valid: 1920×1080 H.264+AAC, \*\*35.09s\*\* duration = exactly the 4 padded clips (8+9+9+9). GCS write access confirmed. Now the \*\*full real-GCS round-trip\*\* ? worker compile fetching a gs:// source + gs:// spec and uploading the reel + marker to gs:// (the exact plumbing the L4

uses):

#### 249. user (2026-07-08T05:59:13.827Z)

=== stage spec to GCS ===

.....  
staged: gs://khelsutra-rally-serving/\_pr525test/\_compile/tok/spec.json (locale=hi ? Hindi  
watermark + Devanagari font)

#### 250. user (2026-07-08T05:59:52.223Z)

=== RUN worker compile with REAL GCS I/O (fetch gs source + spec, upload reel + marker to gs)  
===

C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\site-  
packages\google\auth\\_default.py:113: UserWarning: Your application has authenticated using end  
user credentials from Google Cloud SDK without a quota project. You might receive a "quota  
exceeded" or "API not enabled" error. See the following page for troubleshooting:  
<https://cloud.google.com/docs/authentication/adc-troubleshooting/user-creds>.  
C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\site-  
packages\google\auth\\_default.py:113: UserWarning: Your application has authenticated using end  
user credentials from Google Cloud SDK without a quota project. You might receive a "quota  
exceeded" or "API not enabled" error. See the following page for troubleshooting:  
<https://cloud.google.com/docs/authentication/adc-troubleshooting/user-creds>.  
11:29:38 [INFO] backend.pipeline.stitcher.compiler: Source video duration: 180.05s  
11:29:38 [INFO] backend.pipeline.stitcher.compiler: Trimming 3 segments (begin\_padding=0.5s,  
end\_padding=0.5s)...  
11:29:38 [INFO] backend.pipeline.stitcher.compiler: [watermark] applying branding overlay to  
each clip.  
11:29:38 [WARNING] backend.pipeline.stitcher.compiler: [compile] encoder 'h264\_nvenc' is not  
runnable in this ffmpeg build ? falling back to libx264 (CPU). On the L4 worker this should be  
available.  
11:29:38 [INFO] backend.pipeline.stitcher.compiler: Trimming Segment #1/3 [5.00s - 12.00s] ->  
[4.500s - 12.500s] (duration: 8.00s)  
11:29:39 [INFO] backend.pipeline.stitcher.compiler: Trimming Segment #2/3 [40.00s - 48.00s] ->  
[39.500s - 48.500s] (duration: 9.00s)  
11:29:41 [INFO] backend.pipeline.stitcher.compiler: Trimming Segment #3/3 [90.00s - 98.00s] ->  
[89.500s - 98.500s] (duration: 9.00s)  
11:29:43 [INFO] backend.pipeline.stitcher.compiler: Successfully trimmed 3/3 clips. Preparing to  
stitch...  
11:29:43 [INFO] backend.pipeline.stitcher.compiler: Stitching clips together...  
11:29:43 [INFO] backend.pipeline.stitcher.compiler: Dynamic highlight reel saved:  
C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-  
highlight-  
indexer/75d3d121-f240-4293-afa5-b436a379ad31/scratchpad/compiletest/gcs\_wscratch\smoke-reel.mp4  
(16.2 MB)  
C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\site-  
packages\google\auth\\_default.py:113: UserWarning: Your application has authenticated using end  
user credentials from Google Cloud SDK without a quota project. You might receive a "quota  
exceeded" or "API not enabled" error. See the following page for troubleshooting:  
<https://cloud.google.com/docs/authentication/adc-troubleshooting/user-creds>.  
C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\site-  
packages\google\auth\\_default.py:113: UserWarning: Your application has authenticated using end  
user credentials from Google Cloud SDK without a quota project. You might receive a "quota  
exceeded" or "API not enabled" error. See the following page for troubleshooting:  
<https://cloud.google.com/docs/authentication/adc-troubleshooting/user-creds>.  
11:29:51 [INFO] backend.worker: [worker] wrote compile marker -> gs://khelsutra-rally-  
serving/\_pr525test/\_compile/tok.done  
[worker] compile: video\_id=05e0e577029daafc425ca70c0c4f0a3 clips=3 encoder=h264\_nvenc  
out=smoke-reel.mp4  
[worker] compile: stitching  
[worker] compile: uploading reel  
[worker] compile: finished  
worker-rc=0  
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

### 251. assistant (2026-07-08T06:00:00.729Z)

Real-GCS round-trip succeeded (fetch gs source + spec ? stitch ? **uploading reel** ? marker to gs://, RC=0). Let me verify the reel + marker landed in GCS and the reel is valid, then confirm the marker now carries a populated `reel\_uri`:

### 252. user (2026-07-08T06:00:41.719Z)

```
=== reel mirrored to GCS? ===
 16940481 2026-07-08T05:59:48Z gs://khelsutra-rally-
serving/_pr525test/highlights/05e0e577029daafc425ca70c0c4f0a3/smoke-reel.mp4
TOTAL: 1 objects, 16940481 bytes (16.16MiB)
=== marker in GCS (reel_uri now populated) ===
{"video_id": "05e0e577029daafc425ca70c0c4f0a3", "returncode": 0, "status": "success",
"download_url": "/api/highlights/05e0e577029daafc425ca70c0c4f0a3/smoke-reel.mp4", "reel_uri":
"gs://khelsutra-rally-serving/_pr525test/highlights/05e0e577029daafc425ca70c0c4f0a3/smoke-
reel.mp4", "message": ""}
=== download the GCS reel + ffprobe (real, playable) ===
Average throughput: 35.6MiB/s
codec_name=h264
width=1920
height=1080
codec_name=aac
codec_name=unknown
duration=26.080000
size=16940481
```

### 253. assistant (2026-07-08T06:01:30.508Z)

Fully verified against real GCS: reel mirrored (16.16 MB), marker carries a populated `reel\_uri` + `download\_url` (exactly what the CP reads back), and the downloaded reel is valid (1920×1080, 26.08s = 3 padded clips). The **entire worker path runs correctly against real GCS**; only NVENC + the CP-initiated L4 dispatch remain.

Since ultracode is on, let me launch an adversarial multi-agent review of my diff (to catch bugs before the owner) while I run the final integration test. First, capture the diff for the reviewers:

### 254. user (2026-07-08T06:01:33.798Z)

```
79 C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-
highlight-indexer/75d3d121-f240-4293-afa5-b436a379ad31/scratchpad/review/pr525.diff
saved diff for reviewers
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 255. assistant (2026-07-08T06:01:47.397Z)

Only 79 lines ? that's wrong for a ~1700-line change. Let me diagnose:

### 256. user (2026-07-08T06:01:49.275Z)

```
=== refs ===
dac6736cb23a681135b707d768014820b239a981
8866ffe2757e8fb87b84dcdb4c4158fa1f42f98f
=== stat ===
..._cloud-serving_validation-latency_2026-07-07.md | 48 +++-----
 1 file changed, 6 insertions(+), 42 deletions(-)
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 257. assistant (2026-07-08T06:02:09.523Z)

`origin/master` **moved** (699d85b ? dac6736) ? a sibling session merged the \$9 runlog update while I worked (exactly the "master moves under you" case the safety rules warn about). My branch doesn't touch that file, so the PR stays clean/mergeable. Let me verify branch

cleanliness and regenerate the diff against the true merge-base:

#### 258. user (2026-07-08T06:02:14.378Z)

```
merge-base=699d85bc0abd6e35d3196a4566dcec23f649394e
=== my commits since merge-base (must be ONLY mine) ===
8866ffe perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521)
=== does new master's $9 change conflict with my branch? (files I touch vs master's change) ===
backend/compute/__init__.py
backend/compute/base.py
backend/compute/cloud_run.py
backend/config/__init__.py
backend/config/models.py
backend/config/placement.py
backend/config/presets.py
backend/orchestration.py
backend/pipeline/stitcher/compiler.py
backend/worker/__main__.py
docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md
docs/COMPUTE_DECOUPLED_SERVING/README.md
docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-serving_validation-
latency_2026-07-07.md
tests/test_api_cloud_dispatch.py
tests/test_compile_dispatch.py
tests/test_compiler.py
tests/test_compute_backend.py
tests/test_config_deployment.py
tests/test_phase_placement.py
tests/test_worker.py
=== regenerate diff from merge-base ===
2166 C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-guru-badminton-
highlight-indexer/75d3d121-f240-4293-afa5-b436a379ad31/scratchpad/review/pr525.diff
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

#### 259. assistant (2026-07-08T06:02:36.533Z)

The name-only list shows my backend files in `merge-base..origin/master` ? that suggests \*\*PR #525 may have already been merged\*\* (as a squash). Let me confirm the master state and PR status before anything else:

#### 260. user (2026-07-08T06:02:40.481Z)

```
=== PR 525 state ===
{"mergeCommit":"dac6736cb23a681135b707d768014820b239a981","mergedAt":"2026-07-
08T05:59:33Z","state":"MERGED"}
=== recent origin/master commits ===
dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521) (#525)
bf97700 docs(runlogs): record live SPA CUJ closing the 2 owner-gated validation-latency cells
($9) (#522)
699d85b feat(validation): cache validation verdict by video_id to skip re-validation on reupload
(#518)
71c88ef docs(runlogs): validation-latency closure report for #512 + #513 (cloud-serving) (#519)
115b83d docs(runlogs): record live smoke of the crash-before-marker fast-fail (#515) (#517)
ef41473 fix(cloud-run): fail fast when a worker execution terminates without a done-marker
(#515)
=== dac6736 summary ===
dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521) (#525)
 backend/compute/__init__.py | 2 +
 backend/compute/base.py | 34 +++
 backend/compute/cloud_run.py | 100 ++++++-
 backend/config/__init__.py | 16 ++
```



```

backend/config/models.py | 52 ++++
backend/config/placement.py | 59 ++++
backend/config/presets.py | 10 +-
backend/orchestration.py | 322 ++++++-----
backend/pipeline/stitcher/compiler.py | 94 +++++-
backend/worker/__main__.py | 139 ++++++
docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md | 169 ++++++
docs/COMPUTE_DECOUPLED_SERVING/README.md | 1 +
tests/test_api_cloud_dispatch.py | 70 ++++
tests/test_compile_dispatch.py | 204 ++++++
tests/test_compiler.py | 64 ++++
tests/test_compute_backend.py | 111 ++++++
tests/test_config_deployment.py | 5 +
tests/test_phase_placement.py | 138 ++++++
tests/test_worker.py | 175 ++++++
19 files changed, 1710 insertions(+), 55 deletions(-)
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

#### 261. assistant (2026-07-08T06:03:16.007Z)

**\*\*PR #525 is already MERGED\*\*** (squash `dac6736`, at 05:59 UTC ? merged while I was running the real-video tests). So master now carries all 19 files / 1710 insertions of #521, and my real-video tests above validated the \*merged\* code. The \$9 runlog (#522) is also on master now, so my doc reference resolves.

Since it's merged, any issue a review surfaces should go in a follow-up PR. Let me run a rigorous adversarial multi-agent review of the merged diff (ultracode) to make "thorough testing" real ? findings verified by independent skeptics before I trust them:

#### 262. user (2026-07-08T06:03:19.869Z)

Your configuration specifies to merge with the ref 'refs/heads/feat/521-compile-on-l4' from the remote, but no such ref was fetched.  
 --- worktree still on my branch (has the exact merged code) ---  
 8866ffe perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine placement (#521)  
 Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

#### 263. assistant (2026-07-08T06:03:25.491Z)

Branch was auto-deleted post-merge (normal for squash) ? my worktree at `8866ffe` holds the exact merged code. Now the adversarial review workflow:

#### 264. user (2026-07-08T06:04:07.681Z)

Workflow launched in background. Task ID: ws7lmnehf  
 Summary: Adversarially review the merged #521 compile-on-L4 diff; verify each finding before trusting it  
 Transcript dir: C:\Users\avidu\claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\75d3d121-f240-4293-afa5-b436a379ad31\subagents\workflows\wf\_5f13fc79-bad  
 Script file: C:\Users\avidu\claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\75d3d121-f240-4293-afa5-b436a379ad31\workflows\scripts\pr525-adversarial-review-wf\_5f13fc79-bad.js  
 (Edit this file with Write/Edit and re-invoke Workflow with {scriptPath: "C:\Users\avidu\claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\75d3d121-f240-4293-afa5-b436a379ad31\workflows\scripts\pr525-adversarial-review-wf\_5f13fc79-bad.js"} to iterate without resending the script.)  
 Run ID: wf\_5f13fc79-bad  
 To resume after editing the script: Workflow({scriptPath: "C:\Users\avidu\claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\75d3d121-f240-4293-afa5-b436a379ad31\workflows\scripts\pr525-adversarial-review-wf\_5f13fc79-bad.js", resumeFromRunId: "wf\_5f13fc79-bad"}) ? completed agents return cached results (cached results may themselves be empty ? inspect journal.jsonl before assuming there is

something to recover).

You will be notified when it completes. Use /workflows to watch live progress.

#### 265. assistant (2026-07-08T06:04:34.490Z)

The review runs in the background. Meanwhile, let me update memory (PR merged + tests done) and clean up the GCS test artifacts (leave-no-trace):

#### 266. user (2026-07-08T06:04:42.657Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\pr-525-compile-on-l4.md has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 267. user (2026-07-08T06:05:04.845Z)

```
=== clean up GCS test artifacts (leave-no-trace) ===
Removing gs://khelsutra-rally-serving/_pr525test/_compile/tok/spec.json#1783490354082363...
Removing gs://khelsutra-rally-serving/_pr525test/probe.txt#1783490326257775...
.....
=== verify gone ===
ERROR: (gcloud.storage.ls) One or more URLs matched no objects.
```

#### 268. assistant (2026-07-08T06:05:16.843Z)

GCS artifacts cleaned. Now let me gather the precise deployed-infra facts so I can give the user an accurate L4-test plan (does the control-plane also share `:latest`? what GPU/mounts does the Job need?):

#### 269. user (2026-07-08T06:05:27.443Z)

```
=== control-plane SERVICE image (does it also use :latest?) ===
asia-southeast1-docker.pkg.dev/gen-lang-client-0306026826/rally/rally-worker:latest
=== rally-worker JOB: gpu / mounts / timeout / SA ===
spec:
 template:
 spec:
 template:
 spec:
 containers:
 - resources:
 limits:
 cpu: '8'
 memory: 32Gi
 nvidia.com/gpu: '1'
 volumeMounts:
 - mountPath: /gcs
 name: serving
 nodeSelector:
 run.googleapis.com/accelerator: nvidia-l4
 serviceName: 492532266377-compute@developer.gserviceaccount.com
 timeoutSeconds: '3600'
 volumes:
 - csi:
 driver: gcsfuse.run.googleapis.com
 volumeAttributes:
 bucketName: khelsutra-rally-serving
 name: serving
```

#### 270. assistant (2026-07-08T06:06:37.952Z)

The key infra fact for the L4 test: **both the live control-plane and the worker Job run `rally-worker:latest`** (8 vCPU / L4 / gcsfuse mount of `khelsutra-rally-serving`). The deployed image

predates my `compile` subcommand, so a real L4 exec needs a fresh image build ? which touches the shared live tag. That's a spend + live-infra decision (D17 requires confirming spend first), so I'll tee it up for you rather than deploy unilaterally.

Here's where things stand while the adversarial review finishes:

**\*\*? PR ? done and MERGED.\*\*** #525 squash-merged to master as `dac6736` (all 11 CI checks green). "Open a PR" is satisfied; no duplicate needed.

**\*\*? Real-video testing (on the merged code, no deploy required):\*\*** I ran the actual `python -m backend.worker compile` handler end-to-end on a real 1080p/180s corpus clip:  
- **\*\*Local:\*\*** requested `h264\_nvenc` ? the probe correctly found it unrunnable on this AMD box ? **\*\*fell back to libx264\*\***, trimmed 4 real clips with the branding watermark, produced a valid 35.09s reel. The NVENC-fallback path works on real ffmpeg.  
- **\*\*Real GCS round-trip:\*\*** fetched a `gs://` source + `gs://` spec, stitched, **\*\*uploaded the reel to `gs://khelsutra-rally-serving/?/highlights/<id>/` and wrote the `gs://` done-marker\*\*** (with `reel\_uri` + `download\_url` populated ? exactly what the control-plane reads back). Downloaded reel verified valid. This exercises the *\*exact\** worker code path the L4 runs ? only NVENC hardware encode and the CP-initiated Cloud Run dispatch differ (both proven with mocks). Test artifacts cleaned up from GCS.

**\*\*? Running now:\*\*** an adversarial multi-agent review of the merged diff (5 review dimensions ? each finding refuted-or-confirmed by 2 independent skeptics) to catch anything the mocked tests hide. I'll report verified findings when it lands, and file a follow-up PR if anything real surfaces.

**\*\*? Pending your call ? the real L4 execution.\*\*** It's possible but needs an image build (~8?10 min, ~\$0.5?1 total incl. one L4 exec, well under the D17 cap) and touches the shared `:latest`. I'll lay out the exact options (rebuild `:latest` vs. an isolated throwaway `rally-worker-pr525` Job) with the review results so you can decide with full information.

#### 271. user (2026-07-08T06:13:32.137Z)

<system-reminder>

The user started your suggested background task task\_217edf9c ("Cheaper reel watermark (avoid full re-encode)") in a separate local session. It is running independently. You will be notified here when it ends.

</system-reminder>

dont run the real test. I will make a release today and fold all these in

#### 272. assistant (2026-07-08T06:13:50.121Z)

Understood ? I won't touch the L4 / any deploy. Since you're releasing today, let me make sure the adversarial review lands so any real finding can go into that release. Let me check its status:

#### 273. user (2026-07-08T06:13:53.700Z)

=== dir ===

total 2533

drwxr-xr-x	1	avidu	197609		0 Jul	8 11:42	.
drwxr-xr-x	1	avidu	197609		0 Jul	8 11:34	..
-rw-r--r--	2	avidu	197609	392347	Jul	8 11:41	agent-a0baf7ebe53c21630.jsonl
-rw-r--r--	1	avidu	197609	48	Jul	8 11:34	agent-a0baf7ebe53c21630.meta.json
-rw-r--r--	2	avidu	197609	109226	Jul	8 11:43	agent-a26929d4c969fb9c5.jsonl
-rw-r--r--	2	avidu	197609	437700	Jul	8 11:43	agent-a52a7d189c652e8cf.jsonl
-rw-r--r--	1	avidu	197609	48	Jul	8 11:34	agent-a52a7d189c652e8cf.meta.json
-rw-r--r--	2	avidu	197609	415462	Jul	8 11:43	agent-a588aba9828101736.jsonl
-rw-r--r--	1	avidu	197609	48	Jul	8 11:34	agent-a588aba9828101736.meta.json
-rw-r--r--	2	avidu	197609	123009	Jul	8 11:43	agent-a71a55a475829e001.jsonl
-rw-r--r--	2	avidu	197609	143998	Jul	8 11:43	agent-a77adcbc098c7cca7.jsonl
-rw-r--r--	2	avidu	197609	372457	Jul	8 11:42	agent-a7e807a82054c4b5a.jsonl
-rw-r--r--	1	avidu	197609	48	Jul	8 11:34	agent-a7e807a82054c4b5a.meta.json

```

-rw-r--r-- 2 avidu 197609 94686 Jul 8 11:43 agent-a9ecb0b2f9a51ffea.jsonl
-rw-r--r-- 2 avidu 197609 468507 Jul 8 11:43 agent-af2e3f7269e288069.jsonl
-rw-r--r-- 1 avidu 197609 48 Jul 8 11:34 agent-af2e3f7269e288069.meta.json
-rw-r--r-- 2 avidu 197609 5495 Jul 8 11:42 journal.jsonl
=== journal tail (if present) ===
{"type":"started","key":"v2:e0ded2cd4e19dcc6fbbf04e22dc3900edf2d93a82cb14a97db37b6b966d5a367","a
gentId":"a7e807a82054c4b5a"}
{"type":"started","key":"v2:b51d73523cb7de4cba49008d6e0dc00b811137c64c4aa7ebf0cdf707b3d9e5a9","a
gentId":"a0baf7ebe53c21630"}
{"type":"started","key":"v2:f1a3dd4c46f003f0df7ee70898d1fc42eb8c5ea73cac6636837c9160023b0039","a
gentId":"a588aba9828101736"}
{"type":"started","key":"v2:0f8b659245135fe7e2ec10f1b5c8f40946026825bbcb1040cf1b527e99a55c791","a
gentId":"af2e3f7269e288069"}
{"type":"started","key":"v2:7abb01a326ea80c73d898d3ccf7b6af41ab29be9762b23572bd7c03b865f5f0f","a
gentId":"a52a7d189c652e8cf"}
{"type":"result","key":"v2:b51d73523cb7de4cba49008d6e0dc00b811137c64c4aa7ebf0cdf707b3d9e5a9","ag
entId":"a0baf7ebe53c21630","result":{"findings":[]}}
{"type":"result","key":"v2:e0ded2cd4e19dcc6fbbf04e22dc3900edf2d93a82cb14a97db37b6b966d5a367","ag
entId":"a7e807a82054c4b5a","result":{"findings":[{"file":"backend/orchestration.py","line":689,"
severity":"medium","category":"robustness","summary":"When validation is delegated to the L4
worker (phase_placement.validation=cloud_run_l4), a worker validation REJECTION surfaces to the
user as an opaque \"cloud job failed (returncode 2).\" with no reason, silently dropping the
validation feedback the in-process path provides.\"","failure_scenario":"Config:
deployment.phase_placement.validation=cloud_run_l4 with segment=cloud_run_l4 (or
compute_target=cloud_run). A user uploads a too-blurry / not-badminton clip. _execute_processing
sets validate_on_control_plane=False (orch:865-868), so the control-plane skips its validators
and dispatches with skip_validation=worker_skips_validation=False (orch:666-674). The worker's
blur/relevance validator rejects the clip: cmd_infer hits `if failures: return _fail(2)`
(worker/__main__.py:333-338), which writes a done-marker via _write_done_marker containing ONLY
{video_id, returncode:2, segment_count:0, status:\"failed\"} ? NO message/validator-reason field
(worker/__main__.py:228-238). run_infer returns ComputeResult(returncode=2);
_maybe_dispatch_remote line 688-689 then returns _failed(f\"cloud job failed (returncode
{result.returncode}).\") ? it never reads result.report.get('message') (unlike the compile path
at orch:1393). The user sees \"cloud job failed (returncode 2).\" instead of the specific 'video
too blurry / not badminton' message the in-process control-plane path returns. This failure mode
is newly reachable in #521 (pre-#521 skip_validation was hardcoded True so the worker never
validated) and is untested ? the new tests
(test_validation_delegated_to_worker_when_placed_on_l4 / _not_delegated...) only cover the PASS
path and never assert the worker-rejection message.\"","suggested_fix":"Carry the validation-
failure reason through the delegated path: (1) include a `message` (joined failing-validator
messages) in cmd_infer's failure done-marker (extend _write_done_marker like the compile
marker), and (2) in _maybe_dispatch_remote line 688-689 surface it, e.g. `msg =
str(result.report.get('message')) or f\"cloud job failed (returncode {result.returncode}).\")` ?
mirroring _dispatch_compile_remote at orch:1391-1396. Add a test that a worker validation
rejection returns a reason, not just an rc.\"},{\"file\":\"backend/worker/__main__.py\",\"line\":488,\"s
everity\":\"low\",\"category\":\"robustness\",\"summary\":\"cmd_compile writes its done-marker ONLY after
_stitch_reel returns; any failure BEFORE that (spec fetch/parse, time_pairs/stitching coercion,
VideoCompiler init) raises out of an unwrapped main(), so the worker exits with a traceback and
NO marker ? unlike cmd_infer's _fail() which always writes a marker for every worker-side
failure.\"","failure_scenario":"The bucket-staged spec is truncated/corrupt (transient partial
read) or time_pairs contains a malformed pair. `_read_compile_spec` (worker:426-432) raises
JSONDecodeError/ValueError, or the coercion `[(float(pair[0]), float(pair[1])) for pair ...]`
(worker:493-495) / `StitchingConfig(**...)` (worker:496) raises IndexError/TypeError.
cmd_compile has no surrounding try/except and main() at line 973 calls args.func(args)
unwrapped, so the process exits non-zero with a raw traceback and never writes the --done-uri
marker. The dispatcher's bucket-poll then depends ENTIRELY on the #515 crash-fast Executions-API
probe to convert this into a clean failed job; if that probe is unavailable/delayed, the poll
hangs up to deployment.remote_poll_max_wait_sec (~65 min). cmd_infer avoids this asymmetry ? its
_fail() writes a real-rc marker for every pre-segmentation failure so the dispatcher terminates
fast regardless of the Executions API.\"","suggested_fix":"Give cmd_compile the same discipline as
cmd_infer: wrap spec-read + coercion in try/except and, on failure, write a failed done-marker
(rc + message) via _write_compile_marker before returning a non-zero rc, so a corrupt/partial
spec is a clean fast failed-job rather than relying solely on the crash-fast probe.\"}}]}
{"type":"started","key":"v2:f6f1728ecec9b7e3222de5ecf00b5f86a4e7cd1bfce22e89aa41434a766b8cba","a
gentId":"a9ecb0b2f9a51ffea"}

```

```
{"type":"started","key":"v2:943a9bc5cadee38ce2782798aeb82f817f8d1e098c5c89e63eb92592e8250e2c","agentId":"a26929d4c969fb9c5"}
{"type":"started","key":"v2:0699a9041ce9cddb51ac5784b5b77fd16b2f7170ed8df027f185cbc5c74587ae","agentId":"a71a55a475829e001"}
{"type":"started","key":"v2:acd6fa2cf2ea8decba1a545987f6f305435dd138dbae56d856b9447ed55af7798","agentId":"a77adcbc098c7cca7"}
```

#### 274. assistant (2026-07-08T06:14:30.927Z)

Got it ? **no** real L4 test, no deploy. **I'll** leave `rally-worker` / `:latest` untouched; you fold everything into today's release.

The adversarial review is on its verification pass and has already surfaced **two** real robustness gaps worth folding into that release (both small, both in the newly-added #521 code):

1. **[medium]** Opaque validator-rejection message when validation is delegated to the L4 (`orchestration.py`~689`). With `phase_placement.validation=cloud_run_l4``, a worker rejection surfaces as `"cloud job failed (returncode 2)."` instead of the specific "too blurry / not badminton" reason ? because the `infer` done-marker` carries no `message`` and `_maybe_dispatch_remote`` doesn't read `result.report.get('message')`` (the compile path already does). This is a UX regression for the delegation feature (I'd flagged the risk earlier; the review confirms it's worth fixing).

2. **[low]** `cmd_compile`` has no marker-on-early-failure (`worker/`__main__.py``). A corrupt/partial staged spec raises out of unwrapped `main()`` ? no done-marker ? the dispatcher falls back solely to the #515 crash-fast probe instead of terminating cleanly. `cmd_infer`s _fail()`` always writes a marker; `cmd_compile`` should match.

Let me wait for the verification pass to finish (it may confirm/refute these and add findings from the other dimensions), then I'll fix the **confirmed** set and put up a small follow-up PR you can merge into today's release. I'll report the full verified list shortly.

#### 275. user (2026-07-08T06:20:03.628Z)

```
<task-notification>
<task-id>ws7lmnehf</task-id>
<tool-use-id>toolu_01TjmDrzJMeY9Cv4QhEyfcS3</tool-use-id>
<output-file>C:\Users\avidu\AppData\Local\Temp\claude\C--Users\avidu-Projects-khelsutra-guru-badminton-highlight-indexer\75d3d121-f240-4293-afa5-b436a379ad31\tasks\ws7lmnehf.output</output-file>
<status>completed</status>
<summary>Dynamic workflow "Adversarially review the merged #521 compile-on-L4 diff; verify each finding before trusting it" completed</summary>
<result>{"confirmed":[{"file":"backend/orchestration.py","line":1263,"severity":"high","category":"correctness","summary":"On the L4 worker path, a failed (or skipped) bucket mirror-upload is swallowed and _stitch_reel still returns status='success', so a compile job reports success while the reel is never delivered ? the worker's local file is discarded on Job exit and the control-plane has no local copy to fall back to, so the SPA's download 404s (silent data loss reported as success).","failure_scenario":"phase_placement.compile=cloud_run_l4. Worker cmd_compile -> _stitch_reel stitches the reel into scratch, then storage.put(out_path, reel_uri) raises (transient GCS 5xx / IAM drift / quota) OR reel_uri is empty because the forwarded storage.output_root has no '://' scheme. The except at line 1263 logs a warning and control falls through to `return {'status':'success', ..., 'reel_uri': reel_uri or ''}`. cmd_compile then writes a done-marker with returncode=0/status=success/download_url set. run_compile reads rc=0 and _dispatch_compile_remote returns status='success' with download_url=/api/highlights/&vid&name&name&name. The user clicks it: download_highlight() calls signed_reel_url -> storage.exists(reel_uri)=False (mirror failed) -> None, then falls back to os.path.join(output_dir, name) on the CONTROL-PLANE, where the file does not exist (it was stitched on the worker) -> HTTP 404 'No compiled file named ...'. The reel is lost but the job/SPA both show success.","suggested_fix":"In the worker/no-local-fallback case the bucket mirror is the ONLY delivery mechanism (the code comment at lines 1254-1255 acknowledges this), so a mirror-upload failure ? or an empty reel_uri when a bucket output is required ? must fail the stitch there: return {'status':'failed', 'message': 'reel mirror-upload failed / no bucket destination'} instead of success, so cmd_compile writes rc!=0 and the dispatcher surfaces
```

a clean failed job (letting the user retry). Keep the best-effort swallow only for the true control-plane in-process path where the CP-local file is a valid fallback. Also add an integration test exercising the real `_stitch_reel` on the worker with a failing `storage.put` (current tests in `test_compile_dispatch.py` / `test_worker.py` stub `_stitch_reel`, so this failure mode is uncovered).", "dimension": "correctness", "verdicts": [{"is\_real": true, "confidence": "high", "reasoning": "CONFIRMED in the actual code. On the L4 worker compile path, a failed bucket mirror-upload is swallowed and the job still reports success, producing a download 404 (silent data loss reported as success). Verified chain:\n\n1) backend/orchestration.py:1287-1309 ? after stitching to a scratch out\_path, `_stitch_reel` does `reel_uri = reel_object_uri(...)`; `if reel_uri:` it runs `storage.put(out_path, reel_uri, ...)` in a try/except that on ANY exception only logs `logger.warning("[compile] could not mirror reel to %s...")` (1294-1295) and falls through to `return {"status": "success", ..., "reel_uri": reel_uri or ""}` (1301-1309). The comment at 1282-1286 explicitly acknowledges that on the L4 worker "the mirror is how the reel reaches the bucket at all (the worker's local file is discarded when the Job exits)" ? yet the swallow-and-succeed is applied unconditionally to both the CP caller (where a local fallback exists) and the worker caller (where it does not).\n\n2) backend/worker/\_\_main\_\_.py:506-524 ? `cmd_compile` calls the shared `_stitch_reel`, then `status = result.get("status")` = 'success', `rc = 0` if `status=="success"` else 3 = 0, and `_write_compile_marker` records `status='success' + download_url`.\n\n3) backend/orchestration.py:1435-1468 ? `_dispatch_compile_remote` only checks `if result.returncode != 0:`; on `rc=0` it returns `status='success'` with the `download_url`. It never re-verifies the reel exists in the bucket (no `storage.exists` re-check).\n\n4) backend/main.py:1046-1083 + orchestration.py:1180-1191 ? `download_highlight` calls `signed_reel_url` ? `storage.exists(reel_uri)` is False (mirror failed) ? returns None ? falls back to `os.path.join(output_dir, name)` on the CONTROL-PLANE, where the worker-stitched file never existed ? HTTP 404 "No compiled file named ...".\n\nReachability: the reel put and the marker put are separate operations on very different-sized objects (a large 4K reel vs a tiny JSON marker). A transient/partial failure where the large reel upload fails (timeout/size/quota/mid-transfer blip) while the small marker succeeds is realistic and reachable ? that is precisely the false-success case. A total storage outage would instead also fail the marker put and correctly time the dispatcher out (`rc4/rc5`), so it does not save the partial case. The Cloud Run Job scratch is ephemeral and on a separate instance from the control-plane, so there is genuinely no CP fallback copy.\n\nThe cited line (1263) is off in the current checkout ? the mirror except is at 1294-1295 and the success return at 1301-1309 ? but the described logic is present verbatim and behaves exactly as claimed; the line drift is only a snapshot difference and does not refute the defect. This violates the stated invariant that a stitch/delivery failure must be a clean failed-job.", {"is\_real": true, "confidence": "high", "reasoning": "CONFIRMED ? I could not refute this. The full chain is verified in the merged code.\n\nMechanism (backend/orchestration.py, `_stitch_reel`): after a successful stitch, the reel mirror-upload at lines 1287-1295 is unconditionally best-effort ? `storage.put(out_path, reel_uri, ...)` is wrapped in `except Exception` that only logs `logger.warning("[compile] could not mirror reel to %s...")`. Control then falls through to `return {"status": "success", ..., "reel_uri": reel_uri or ""}` (1301-1309). This helper is SHARED by the control-plane handler and the L4 worker with NO distinction, even though its own comment (1282-1286) states the worker's local file is discarded on Job exit and "the mirror is how the reel reaches the bucket at all." So a swallowed mirror failure that is harmless on the CP (local fallback exists) becomes silent data loss on the worker.\n\nEnd-to-end confirmation:\n- Worker (backend/worker/\_\_main\_\_.py `cmd_compile`, 506-525): `status/rc` are taken directly from the `_stitch_reel` result; on a swallowed mirror failure it gets `status='success'`, `rc=0`, and `_write_compile_marker` (446-461) writes a marker with `rc=0/status='success'/download_url/reel_uri`. No re-verification that the reel actually landed.\n- Dispatcher (backend/compute/cloud\_run.py `_poll_check` 405-407): a PRESENT marker is authoritative ("success signal ? authoritative even if status later flips"), so the crash-fast `RemoteExecutionFailed` path is NOT taken (it requires an ABSENT marker). `run_compile` returns `returncode=0`; `_dispatch_compile_remote` (orchestration.py 1435-1468) returns `status='success'` with `download_url=/api/highlights/<vid>/<name>`.\n- Download (backend/main.py `download_highlight` 1060-1081): `signed_reel_url` ? `reel_object_uri` ? `storage.exists(reel_uri)=False` (upload failed) ? None ? falls back to `os.path.join(output_dir, name)` on the CONTROL-PLANE, where the file never existed (it was stitched on the worker) ? HTTP 404 "No compiled file named '<name>'". Job/SPA both showed success; reel is lost.\n\nRejected refutations: (a) the reel put and the tiny marker put are independent `storage.put` calls ? a transient 5xx / partial-upload / quota / IAM hit on the large reel object while the sub-1KB marker succeeds seconds later is a realistic, non-contrived failure mode (the large object has far more failure surface). (b) The crash-fast path cannot help ? it only fires on an absent marker; here the marker is present and says success. (c) No `verify-reel-exists` step exists anywhere in the compile/worker/dispatch/download path. Both sub-cases in the claim manifest: an empty `reel_uri` (misprovisioned `output_root` without `://`) skips the mirror entirely

```

and st
... (truncated 9852 chars, full result in C:\Users\avidu\AppData\Local\Temp\claude\C--Users-
avidu-Projects-khelsutra-guru-badminton-highlight-
indexer\75d3d121-f240-4293-afa5-b436a379ad31\tasks\ws7lmnehf.output)</result>
<diagnostics>Per-agent results: C:\Users\avidu\claude\projects\C--Users-avidu-Projects-
khelsutra-guru-badminton-highlight-
indexer\75d3d121-f240-4293-afa5-b436a379ad31\subagents\workflows\wf_5f13fc79-bad\journal.jsonl ?
one {"type":"result",...} line per completed agent with its full return value.
If the result above is empty or unexpected, Read this file BEFORE diagnosing ? do not assume
agents returned non-empty results.
To re-run with edited post-processing: Workflow({scriptPath: 'C:\Users\avidu\claude\projects\C
--Users-avidu-Projects-khelsutra-guru-badminton-highlight-
indexer\75d3d121-f240-4293-afa5-b436a379ad31\workflows\scripts\pr525-adversarial-review-
wf_5f13fc79-bad.js', resumeFromRunId: 'wf_5f13fc79-bad', args:
{"diffPath":"C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-khelsutra-
guru-badminton-highlight-
indexer/75d3d121-f240-4293-afa5-b436a379ad31/scratchpad/review/pr525.diff","root":"C:/kg-
wt521\"}"} ? agents whose (prompt, opts) are unchanged replay from cache.</diagnostics>
<usage><agent_count>19</agent_count><agents_done>19</agents_done><agents_error>0</agents_error><
agents_skipped>0</agents_skipped><agents_empty_result>1</agents_empty_result><subagent_tokens>18
23930</subagent_tokens><tool_uses>305</tool_uses><duration_ms>955885</duration_ms></usage>
</task-notification>

```

## 276. assistant (2026-07-08T06:20:17.422Z)

The review completed (19 agents, ~16 min). It found a **HIGH-severity bug I missed**. Let me read the full verified result:

## 277. user (2026-07-08T06:20:26.871Z)

```

Exit code 1
Traceback (most recent call last):
 File "<string>", line 4, in <module>
 c=d['confirmed']
 ~^^^^^^^^^^^^^^^^
KeyError: 'confirmed'

```

## 278. user (2026-07-08T06:20:34.431Z)

```

TOP KEYS: ['summary', 'agentCount', 'logs', 'result', 'workflowProgress', 'totalTokens',
'totalToolCalls']
--- head of file ---
{
 "summary": "Adversarially review the merged #521 compile-on-L4 diff; verify each finding
before trusting it",
 "agentCount": 19,
 "logs": [
 "review complete: 7 raw findings, 2 confirmed, 5 refuted"
],
 "result": {
 "confirmed": [
 {
 "file": "backend/orchestration.py",
 "line": 1263,
 "severity": "high",
 "category": "correctness",
 "summary"

```

## 279. user (2026-07-08T06:20:43.803Z)

CONFIRMED: 2 | refuted: 5

=== HIGH | correctness | backend/orchestration.py : 1263  
SUMMARY: On the L4 worker path, a failed (or skipped) bucket mirror-upload is swallowed and

`_stitch_reel` still returns `status='success'`, so a compile job reports success while the reel is never delivered ? the worker's local file is discarded on Job exit and the control-plane has no local copy to fall back to, so the SPA's download 404s (silent data loss reported as success).  
FIX: In the worker/no-local-fallback case the bucket mirror is the ONLY delivery mechanism (the code comment at lines 1254-1255 acknowledges this), so a mirror-upload failure ? or an empty `reel_uri` when a bucket output is required ? must fail the stitch there: return `{'status':'failed', 'message': 'reel mirror-upload failed / no bucket destination'}` instead of success, so `cmd_compile` writes `rc!=0` and the dispatcher surfaces a clean failed job (letting the user retry). Keep the best-effort swallow only for the true control-plane in-process path where the CP-local file is a valid fallback. Also add an integration test exercising the real `_stitch_reel` on the worker with a failing `storage.put` (current

=== HIGH | test-coverage | backend/orchestration.py : 1263

SUMMARY: No test exercises the L4 worker's real `_stitch_reel` mirror-upload; a swallowed mirror failure there reports success (`rc 0`) but the reel is lost, so the SPA gets a 404 for a job it was told succeeded.

FIX: On the worker path a mirror-upload failure must FAIL the stitch (return `status=failed / non-zero rc`), because unlike the control-plane the worker's local reel does not survive the Job. Distinguish the two callers (e.g. a `'require_mirror'/'is_worker'` flag into `_stitch_reel`), and add a test that runs `cmd_compile` through the real `_stitch_reel` with `storage.put` patched to raise, asserting `rc != 0` and a failed marker (no false success).

--- REFUTED (dropped) ---

- `backend/orchestration.py` : A cloud-dispatched compile with a non-URI `storage.output_root` silently discards the reel yet reports success, because re
- `backend/pipeline/stitcher/compiler.py` : The NVENC availability probe does not use the same encoder args as the real per-clip encode, so it can report an encoder
- `backend/orchestration.py` : When validation is delegated to the L4 worker (`phase_placement.validation=cloud_run_l4`), a worker validation REJECTION s
- `backend/worker/__main__.py` : `cmd_compile` writes its done-marker ONLY after `_stitch_reel` returns; any failure BEFORE that (spec fetch/parse, time\_pair
- `backend/worker/__main__.py` : The dispatcher->worker contract (compile argv flags and the staged spec JSON schema) is never round-tripped end-to-end;

## 280. assistant (2026-07-08T06:21:55.274Z)

The verification pass cut through it: `**2 confirmed` (both the same HIGH-severity bug), 5 refuted.\*\* Notably, the two issues I mentioned in my last message (opaque validation message; `'cmd_compile' no-early-marker`) were `**refuted**` by the skeptics ? the validation one because delegation is opt-in/not preset-enabled, and the marker one because the #515 crash-fast probe already converts a worker crash into a clean `'rc 5'`. Good ? that's exactly why adversarial verification matters.

The `**real**` bug is one I missed:

`**? HIGH ? silent data loss on the L4 worker: a failed mirror-upload reports 'success'.`  
`'_stitch_reel'`'s reel mirror-upload is best-effort (swallow+warn) ? correct for the `*control-plane*` path (the local file is a valid fallback), but `**wrong for the worker**`, whose local file is discarded when the Job exits. So if `'storage.put'` fails (transient GCS 5xx / quota) `*or*` `'reel_uri'` is empty (misprovisioned `'output_root'`), the worker still returns `'status=success'` ? success marker ? the SPA gets a `'download_url'` that `**404s**`. The shared-core refactor introduced this: the mirror was best-effort when only the CP called it.

This is worth fixing before your release. Let me put it in a follow-up PR off the current master. First, a fresh isolated worktree off `'origin/master'`:

## 281. user (2026-07-08T06:22:09.949Z)

```
origin/master=dac6736
branch 'fix/521-worker-mirror-required' set up to track 'origin/master'.
Updating files: 67% (372/552)
Updating files: 68% (376/552)
Updating files: 69% (381/552)
Updating files: 70% (387/552)
Updating files: 71% (392/552)
```



```

Updating files: 72% (398/552)
Updating files: 73% (403/552)
Updating files: 74% (409/552)
Updating files: 75% (414/552)
Updating files: 76% (420/552)
Updating files: 77% (426/552)
Updating files: 78% (431/552)
Updating files: 79% (437/552)
Updating files: 80% (442/552)
Updating files: 81% (448/552)
Updating files: 82% (453/552)
Updating files: 83% (459/552)
Updating files: 84% (464/552)
Updating files: 85% (470/552)
Updating files: 86% (475/552)
Updating files: 87% (481/552)
Updating files: 88% (486/552)
Updating files: 89% (492/552)
Updating files: 90% (497/552)
Updating files: 91% (503/552)
Updating files: 92% (508/552)
Updating files: 93% (514/552)
Updating files: 94% (519/552)
Updating files: 95% (525/552)
Updating files: 96% (530/552)
Updating files: 97% (536/552)
Updating files: 98% (541/552)
Updating files: 99% (547/552)
Updating files: 100% (552/552)
Updating files: 100% (552/552), done.
HEAD is now at dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven
phase?machine placement (#521) (#525)
--- verify ---
fix/521-worker-mirror-required
dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521) (#525)
--- only my commits (empty) ---
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

## 282. assistant (2026-07-08T06:22:17.183Z)

Clean worktree off `dac6736`. Now let me read the merged `\_stitch\_reel` in the new worktree to apply the exact fix:

## 283. user (2026-07-08T06:22:20.966Z)

```

1172: reel_uri = reel_object_uri(cfg, video_id, name)
1190: def _stitch_reel(
1256: reel_uri = reel_object_uri(cfg, video_id, out_name)
1264: logger.warning("[compile] could not mirror reel to %s: %s", reel_uri, e)
1319: return _stitch_reel(
1423: def _execute_compile(cfg, db, job: Dict[str, Any], progress=None) -> Dict[str, Any]:
1458: return _stitch_reel(
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

## 284. user (2026-07-08T06:22:24.438Z)

```

1190 def _stitch_reel(
1191 cfg,
1192 stitching,
1193 video_id: str,
1194 source_ref: str,
1195 time_pairs: List[Tuple[float, float]],
1196 out_name: str,

```

```

1197 locale: Optional[str],
1198 notify,
1199) -> Dict[str, Any]:
1200 """Stitch + mirror-upload core, shared by the control-plane compile handler AND the
1201 L4 worker
1202 (#521). ``time_pairs`` are already RESOLVED (the caller filtered via
1203 ``_resolve_compile``), so this
1204 needs NO control-plane DB ? the worker reads them from the staged spec. Builds the
1205 VideoCompiler
1206 with ``stitching.encoder`` (NVENC on the L4, libx264 on the CP), fetches the source,
1207 runs the
1208 trim/watermark/concat, mirrors the reel to
1209 ``<output_root>/highlights/<video_id>/<name>``, and
1210 returns the ``{status, filepath, download_url, video_id, reel_uri}`` the SPA polls
1211 for. Returns a
1212 ``status='failed'`` result on a stitch error (never raises for an expected ffmpeg
1213 failure)."""
1214 import backend.main as _main # call-time seam: VideoCompiler stays patchable on
1215 backend.main
1216 output_dir = (
1217 getattr(getattr(cfg, "output", None), "output_dir", "output") or "output"
1218)
1219 # Localize the branding watermark to the UI locale the SPA rendered in (recorded on
1220 the job);
1221 # unknown/absent locale keeps the configured default unchanged.
1222 localized_watermark = localize_watermark(stitching.watermark, locale)
1223 compiler = _main.VideoCompiler(
1224 output_dir,
1225 begin_padding=stitching.begin_padding,
1226 end_padding=stitching.end_padding,
1227 watermark=localized_watermark,
1228 encoder=getattr(stitching, "encoder", "libx264"),
1229)
1230 notify("stitching")
1231 local_src = ""
1232 try:
1233 # The stored `filepath` is the SOURCE ref (gs://, s3://, gdrive://, or a local
1234 path). Resolve to
1235 # a LOCAL file the stitcher can read ? identity for a local path, fetch-to-
1236 scratch for a URI (the
1237 # same seam the process path uses); without it, compiling a bucket-ingested
1238 video can't open gs://.
1239 local_src = storage.acquire_local_input(
1240 source_ref, getattr(cfg, "storage", None)
1241)
1242 # The ffmpeg stitch is CPU-heavy ? serialize it so a raised drain worker count
1243 (for remote
1244 # dispatch) doesn't run multiple compiles at once on the box (see
1245 _LOCAL_COMPUTE_LANE).
1246 with _LOCAL_COMPUTE_LANE:
1247 out_path = compiler.compile_highlights(local_src, time_pairs, out_name)
1248 except Exception as e: # noqa: BLE001 ? a stitch/ffmpeg failure is a clean job
1249 FAILURE, not a crash
1250 logger.error(
1251 "[compile] stitching failed for video_id=%s: %s", video_id, e, exc_info=True
1252)
1253 return {
1254 "status": "failed",
1255 "message": f"Highlight stitching failed: {e}",
1256 "video_id": video_id,
1257 }
1258 finally:
1259 if local_src:

```

```

1247 storage.cleanup_acquired_local_input(
1248 source_ref, local_src, getattr(cfg, "storage", None)
1249)
1250
1251 # #463: mirror the reel to the output bucket so /api/highlights can 307 to a signed
1252 URL ? Cloud Run
1253 # caps a buffered response at ~32 MiB (a real reel is bigger) and the local /tmp
1254 copy dies when the
1255 # instance recycles. A mirror-upload failure must NOT fail an otherwise-good stitch
1256 (the endpoint
1257 # still falls back to the local file on the same instance). On the L4 worker the
1258 mirror is how the
1259 # reel reaches the bucket at all (the worker's local file is discarded when the Job
1260 exits).
1261 reel_uri = reel_object_uri(cfg, video_id, out_name)
1262 if reel_uri:
1263 try:
1264 notify("uploading reel")
1265 # Thread the active storage settings so the configured backend (S3
1266 endpoint/region, GCS
1267 # project/client) is used, not a fresh default.
1268 storage.put(out_path, reel_uri, config=_storage_settings(cfg))
1269 except Exception as e: # noqa: BLE001 ? best-effort mirror; local stitch
1270 already succeeded
1271 logger.warning("[compile] could not mirror reel to %s: %s", reel_uri, e)
1272
1273 # `filepath` is server-local (not browser-reachable); also return the download URL
1274 so the SPA needn't
1275 # hand-build it (P6 / PER_RALLY_SELECTION.md §4.3).
1276 download_url = f"/api/highlights/{quote(video_id)}/{quote(out_name)}"
1277 notify("finished")

```

#### 285. assistant (2026-07-08T06:22:29.322Z)

Now apply the fix ? add a `require\_mirror` flag so the worker (no local fallback) fails on a lost mirror instead of falsely succeeding:

#### 286. user (2026-07-08T06:22:40.384Z)

The file C:\kg-wt525fix\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 287. user (2026-07-08T06:22:56.524Z)

The file C:\kg-wt525fix\backend\orchestration.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 288. assistant (2026-07-08T06:22:59.384Z)

Now make the worker's `cmd\_compile` pass `require\_mirror=True`:

#### 289. user (2026-07-08T06:23:02.723Z)

468: the reel via the SAME ``orchestration.\_stitch\_reel`` core the control-plane uses (so there is no

```

503: result = orchestration._stitch_reel(
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

#### 290. user (2026-07-08T06:23:08.421Z)

```

503 result = orchestration._stitch_reel(
504 config,
505 stitching,
506 video_id,

```

```

507 source_ref,
508 time_pairs,
509 out_name,
510 locale,
511 lambda stage: print(f"[worker] compile: {stage}"),
512)
513 status = str(result.get("status") or "failed")
514 rc = 0 if status == "success" else 3
515 if status != "success":
516 print(f"[worker] compile FAILED: {result.get('message')}")
517

```

#### 291. user (2026-07-08T06:23:15.472Z)

The file C:\kg-wt525fix\backend\worker\\_\_main\_\_.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 292. assistant (2026-07-08T06:23:53.279Z)

Now add tests that exercise the **real** `_stitch_reel` (the finding noted the current tests stub it). Let me append `require_mirror` tests to `test_compile_dispatch.py`:

#### 293. user (2026-07-08T06:23:55.787Z)

```

def test_compile_missing_video_fails_before_dispatch(compile_env):
 """A compile for an unknown video_id fails at resolve (404) before any dispatch."""
 db, cfg, _video_id, _seg_ids, _bucket, monkeypatch = compile_env
 fake = _FakeCompileClient()
 _inject_fake(monkeypatch, fake)
 out = orch._execute_compile(cfg, db, _job("nope", [1, 2]))
 assert out["status"] == "failed"
 assert fake.calls == [] # never dispatched ? resolve failed first
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

#### 294. user (2026-07-08T06:23:58.769Z)

```

197 def test_compile_missing_video_fails_before_dispatch(compile_env):
198 """A compile for an unknown video_id fails at resolve (404) before any dispatch."""
199 db, cfg, _video_id, _seg_ids, _bucket, monkeypatch = compile_env
200 fake = _FakeCompileClient()
201 _inject_fake(monkeypatch, fake)
202 out = orch._execute_compile(cfg, db, _job("nope", [1, 2]))
203 assert out["status"] == "failed"
204 assert fake.calls == [] # never dispatched ? resolve failed first

```

#### 295. user (2026-07-08T06:24:19.927Z)

The file C:\kg-wt525fix\tests\test\_compile\_dispatch.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 296. assistant (2026-07-08T06:24:24.456Z)

Now the worker end-to-end test ? `cmd_compile` through the **real** `_stitch_reel` with a failing upload. Let me add it after the existing compile tests:

#### 297. user (2026-07-08T06:24:27.146Z)

```

1089: def test_cmd_compile_no_done_uri_writes_no_marker(tmp_path, monkeypatch):
1090:)
1091: cfg = tmp_path / "c.json"
1092: cfg.write_text({'storage': {'backend': 'local'}})
1093: rc = wmain.cmd_compile(
1094: wmain.build_parser().parse_args(
1095: [

```

```

 "compile",
 "--spec-uri",
 spec_uri,
 "--out-uri",
 str(bucket),
 "--scratch",
 str(tmp_path / "s"),
 "--config",
 str(cfg),
]
)
)
assert rc == 0
assert not (bucket / "_compile" / "tok.done").exists()
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

#### 298. user (2026-07-08T06:24:30.854Z)

```

1108 "--spec-uri",
1109 spec_uri,
1110 "--out-uri",
1111 str(bucket),

```

#### 299. user (2026-07-08T06:24:33.882Z)

```

1115 str(cfg),
1116]
1117)
1118)
1119 assert rc == 0
1120 assert not (bucket / "_compile" / "tok.done").exists()
1121

```

#### 300. user (2026-07-08T06:24:49.509Z)

The file C:\kg-wt525fix\tests\test\_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 301. assistant (2026-07-08T06:25:05.371Z)

The output\_root must be a `gs://` URI (so the mirror is \*attempted\* and then fails) ? with a local path it'd hit the no-bucket branch instead. Let me fix that:

#### 302. user (2026-07-08T06:25:08.738Z)

The file C:\kg-wt525fix\tests\test\_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 303. assistant (2026-07-08T06:25:20.097Z)

Let me run the fix's tests plus the full compile/worker suites:

#### 304. user (2026-07-08T06:25:58.309Z)

```

 config,
 stitching,
 video_id,
 source_ref,
 time_pairs,
 out_name,
 locale,
 lambda stage: print(f"[worker] compile: {stage}"),
 # On the L4 worker the bucket mirror is the ONLY reel delivery ? the worker's local
file is

```

```
 # discarded when the Job exits. A lost mirror MUST fail the stitch (not a
404-on-"success").
 require_mirror=True,
)
E TypeError: test_cmd_compile_stitches_and_writes_marker.<locals>._fake_stitch() got an
unexpected keyword argument 'require_mirror'
```

```
backend\worker__main__.py:503: TypeError
----- Captured stdout call -----
[worker] compile: video_id=vidReel clips=2 encoder=h264_nvenc out=reel.mp4
===== short test summary info =====
FAILED tests/test_worker.py::test_cmd_compile_stitches_and_writes_marker - Ty...
1 failed, 46 passed in 21.51s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 305. assistant (2026-07-08T06:26:09.529Z)

Two stub `\_fake\_stitch` signatures need the new kwarg. Let me update them to accept it:

### 306. user (2026-07-08T06:26:12.284Z)

The file C:\kg-wt525fix\tests\test\_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

### 307. user (2026-07-08T06:26:17.808Z)

```
1041: assert seen["encoder"] == "h264_nvenc"
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 308. user (2026-07-08T06:26:21.152Z)

```
1040 # the shared stitch core got the resolved pairs + the NVENC encoder from the spec
1041 assert seen["encoder"] == "h264_nvenc"
1042 assert seen["pairs"] == [(2.0, 6.0), (10.0, 14.0)]
1043 assert seen["source"] == str(source)
1044 marker = json.loads(open(done, encoding="utf-8").read())
1045 assert marker["returncode"] == 0 and marker["status"] == "success"
```

### 309. user (2026-07-08T06:26:27.607Z)

The file C:\kg-wt525fix\tests\test\_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

### 310. user (2026-07-08T06:26:32.107Z)

The file C:\kg-wt525fix\tests\test\_compile\_dispatch.py has been updated successfully. (file state is current in your context ? no need to Read it back)

### 311. user (2026-07-08T06:26:43.079Z)

```
..... [100%]
47 passed in 5.17s
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 312. assistant (2026-07-08T06:26:48.418Z)

Fix tests pass. Now the full verification gate (ruff, mypy, leak-scan, link-check, then full suite):

### 313. user (2026-07-08T06:28:06.929Z)

```
=== ruff ===
All checks passed!
```

```

=== mypy ===
backend\providers__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 150 source files
=== leak_scan ===
[leak-scan] clean ? no MD5 / personal-path / private-name leaks in the code surface.
=== md_links ===
[check-md-links] clean ? no broken internal links in live docs.
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

#### 314. assistant (2026-07-08T06:28:15.763Z)

Let me add a brief note to the phase-placement doc about the worker's required-mirror invariant, then run the full suite:

#### 315. user (2026-07-08T06:28:17.727Z)

```

91: core** (no forked pipeline), mirror-uploads it to
`<output_root>/highlights/<video_id>/<name>`,
92: and writes the marker (rc + `download_url` + `reel_uri`).
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

#### 316. user (2026-07-08T06:28:21.350Z)

```

89 one bucket never collide on a marker.
90 5. The worker (`cmd_compile`) reads the spec, stitches the reel via the **shared
`_stitch_reel`
91 core** (no forked pipeline), mirror-uploads it to
`<output_root>/highlights/<video_id>/<name>`,
92 and writes the marker (rc + `download_url` + `reel_uri`).
93 6. The control-plane returns `{status, filepath, download_url, video_id}` ? the reel is
already in the
94 bucket, so `/api/highlights` serves it via the existing signed-URL 307. On a worker
rc?0 ?

```

#### 317. user (2026-07-08T06:28:29.001Z)

The file C:\kg-wt525fix\docs\COMPUTE\_DECOUPLED\_SERVING\PHASE\_PLACEMENT.md has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 318. user (2026-07-08T06:37:47.645Z)

```

[check-md-links] clean ? no broken internal links in live docs.
=== FULL SUITE ===
TOTAL 12669 1805 86%
Required test coverage of 70% reached. Total coverage: 85.75%
1879 passed, 8 skipped, 25 warnings in 547.42s (0:09:07)
=====
RESULT: PASS - all tests green
=====
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer

```

#### 319. assistant (2026-07-08T06:38:12.635Z)

Full suite green (\*\*1879 passed\*\*, coverage 85.75%).

I also notice a sibling session (the watermark follow-up task you started) has surfaced a **separate, important bug** in MEMORY.md as PR #527: my NVENC probe uses a **64x64** test frame, which is **below NVENC's minimum resolution** ? so on a real L4 the probe could falsely report `h264_nvenc` unavailable and silently fall back to libx264, **defeating #521's NVENC win entirely**. That's a real bug, being fixed independently in #527 (different file ? `compiler.py::_probe_encoder_available` ? no conflict with my `_stitch_reel` fix). Good catch by that session; I'll leave it to #527 and keep this PR focused on the mirror-data-loss fix.

Let me verify branch cleanliness and commit (fetching first, since master is moving under concurrent sessions):

### 320. user (2026-07-08T06:38:18.244Z)

```
origin/master=dac6736 myHEAD-base=dac6736
=== my branch ===
fix/521-worker-mirror-required
=== does master now overlap my files? (fetch drift check) ===
=== my changed files ===
 M backend/orchestration.py
 M backend/worker/__main__.py
 M docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md
 M tests/test_compile_dispatch.py
 M tests/test_worker.py
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 321. assistant (2026-07-08T06:38:22.109Z)

Clean ? `origin/master` still at `dac6736`, no file overlap. Let me stage explicit paths and commit:

### 322. user (2026-07-08T06:38:36.438Z)

```
[fix/521-worker-mirror-required 7a31a9b] fix(compile): fail the L4-worker stitch when the reel
mirror-upload fails (#521 follow-up)
 5 files changed, 210 insertions(+), 9 deletions(-)
=== verify: only my commit ahead ===
7a31a9b fix(compile): fail the L4-worker stitch when the reel mirror-upload fails (#521 follow-
up)
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 323. user (2026-07-08T06:38:58.748Z)

```
remote:
branch 'fix/521-worker-mirror-required' set up to track 'origin/fix/521-worker-mirror-required'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
 * [new branch] fix/521-worker-mirror-required -> fix/521-worker-mirror-required
https://github.com/Khelsutra/badminton-highlight-indexer/pull/528
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 324. assistant (2026-07-08T06:39:02.524Z)

Fix PR **\*\*#528\*\*** opened. Let me watch its CI to green:

### 325. user (2026-07-08T06:42:39.793Z)

```
Mypy (lenient baseline) pass 36s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28922900489/job/85803837201
Ruff lint pass 8s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28922900489/job/85803837194
Syntax sweep + import smoke (no GPU stack) pass 18s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28922900489/job/85803837207
Built wheel smoke (install + launch) pass 1m11s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28922900489/job/85803837217
Dependency license scan (no copleft in runtime deps) pass 33s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28922900489/job/85803837214
Full test suite (offline, mocked) pass 3m14s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28922900489/job/85803837265
Golden fixtures cross-OS (macos-latest) pass 22s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28922900489/job/85803837235
Golden fixtures cross-OS (ubuntu-latest) pass 21s
```



```
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28922900489/job/85803837234
Golden fixtures cross-OS (windows-latest) pass 1m0s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28922900489/job/85803837211
Leak scan (PII / attribution guard) pass 8s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28922900489/job/85803837218
Markdown link check (live docs) pass 6s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/28922900489/job/85803837191
Mypy (lenient baseline) pass 36s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28922900489/job/85803837201
Ruff lint pass 8s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28922900489/job/85803837194
Syntax sweep + import smoke (no GPU stack) pass 18s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/28922900489/job/85803837207
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 326. assistant (2026-07-08T06:42:49.625Z)

All 11 CI checks green on **\*\*#528\*\***. Let me record it in memory (the sibling session added #527; I'll append #528):

### 327. user (2026-07-08T06:42:50.134Z)

```
1 # Memory index
2
3 - [PR #525 ? compile-on-L4 (#521)](pr-525-compile-on-l4.md) ? config-driven
phase?machine placement + L4 NVENC stitch; Part 3 (cheaper watermark) deferred.
4 - [PR #527 ? NVENC probe 64x64?256x256](pr-527-nvenc-probe-resolution.md) ? sub-minimum
probe silently defeated #521's NVENC-on-L4 win; open, awaits owner live-GPU confirm.
5
```

### 328. user (2026-07-08T06:42:57.066Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

### 329. assistant (2026-07-08T06:43:09.280Z)

```

name: pr-528-worker-mirror-required
description: PR #528 (#521 follow-up) ? L4-worker compile fails on a lost reel mirror-upload
instead of a silent 404-on-success
metadata:
 type: project

```

**\*\*PR [#528](https://github.com/Khelsutra/badminton-highlight-indexer/pull/528)\*\*** ? open 2026-07-08, all 11 CI checks green, for today's release.

Fixes a HIGH-severity silent-data-loss bug an **\*\*adversarial multi-agent review\*\*** of the merged #521 ([pr-525-compile-on-l4]) surfaced + verified: ``orchestration._stitch_reel``'s reel mirror-upload was best-effort (swallow+warn) ? correct for the control-plane (CP-local reel is a fallback), but WRONG on the L4 worker (its local reel is discarded when the Job exits). A failed ``storage.put`` or an empty ``reel_uri`` (non-bucket ``output_root``) let the worker return ``status=success`` ? success marker ? a ``download_url`` that 404s.

**\*\*Fix:\*\*** ``_stitch_reel(..., require_mirror=False)``; ``worker.cmd_compile`` passes ``require_mirror=True`` ? a lost mirror / no-bucket-destination returns ``status=failed`` (rc 3) ? clean retryable failed job. CP path unchanged (default False).

**\*\*Review verdict (why trust it):\*\*** 7 raw findings ? 2 confirmed (this bug, ×2 dimensions), 5

refuted by independent skeptics ? incl. the opaque-delegated-validation-message (opt-in, not preset-enabled) and cmd\_compile-no-early-marker (#515 crash-fast probe covers it). So those were deliberately NOT changed.

**\*\*Sibling concurrent fix:\*\*** [[pr-527-nvenc-probe-resolution]] fixes the NVENC probe's 64x64 sub-minimum test frame (would silently defeat NVENC on the real L4) ? different file, no conflict. Both fold into today's release alongside #525.

### 330. user (2026-07-08T06:43:09.323Z)

File created successfully at: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\pr-528-worker-mirror-required.md (file state is current in your context ? no need to Read it back)

### 331. assistant (2026-07-08T06:43:38.325Z)

Done. Here's the complete picture for your release today.

## What I did (this session)

**\*\*Real-video testing ? done, no deploy (as you instructed).\*\*** I did **\*not\*** touch `rally-worker`, `:latest`, or run any GPU job. Instead I exercised the actual worker `compile` path on a real 1080p/180s corpus clip:

- **\*\*Real ffmpeg:\*\*** requested `h264\_nvenc` ? probe correctly found it unrunnable on this box ? fell back to libx264, trimmed real clips + watermark ? valid reel.
- **\*\*Real GCS round-trip:\*\*** fetched a `gs://` source + `gs://` spec, stitched, uploaded the reel + wrote the done-marker to `gs://khelsutra-rally-serving/?` (marker carried `reel\_uri` + `download\_url`). Verified the reel, then cleaned up the GCS scratch. This is the exact worker code path the L4 runs ? only NVENC hardware encode and the CP-initiated dispatch differ (both covered by tests).

**\*\*Adversarial review caught a real bug I'd missed.\*\*** A 19-agent verify-before-trust review (5 dimensions ? each finding refuted-or-confirmed by 2 skeptics) returned **\*\*2 confirmed / 5 refuted\*\***:

- ? **\*\*HIGH (fixed):\*\*** on the L4 worker a **\*failed reel mirror-upload\*** returned `status=success` ? the SPA gets a `download\_url` that **\*\*404s\*\***. The worker has no local fallback (its file dies with the Job), so a lost mirror is a lost reel ? but was reported as success. Fixed in **\*\*PR #528\*\*** (`require\_mirror=True` on the worker path ? clean retryable failure; control-plane path unchanged).
- The two issues I'd flagged earlier (opaque delegated-validation message; `cmd\_compile` no-early-marker) were **\*\*refuted\*\*** on verification ? opt-in feature / already covered by the #515 crash-fast probe ? so I deliberately left them.

## PRs for today's release

PR	State	What
[#525](https://github.com/Khelsutra/badminton-highlight-indexer/pull/525)	merged	(`dac6736`)   #521 ? config-driven placement + compile-on-L4 + NVENC
[#528](https://github.com/Khelsutra/badminton-highlight-indexer/pull/528)	open, CI green	the mirror-data-loss fix above
[#527](https://github.com/Khelsutra/badminton-highlight-indexer/pull/527)	open (sibling session)	NVENC probe 64x64?256x256 ? the 64x64 frame is <b>**below NVENC's minimum resolution**</b> , which would silently defeat NVENC on the real L4. Independent file, no conflict with #528.

**\*\*Worth knowing before you cut the release:\*\*** #527 is not cosmetic ? without it, #521's headline NVENC win could silently no-op on the L4 (probe fails on the tiny frame ? libx264 fallback). I'd fold **\*\*#528 and #527\*\*** in alongside #525. Both are green/small.

All checks on #528: full suite **\*\*1879 passed\*\***, ruff + mypy + leak-scan + link-check clean, coverage 85.8%. Isolated worktrees (`C:/kg-wt521`, `C:/kg-wt525fix`) ? I never touched the shared checkout or sibling branches.

### 332. user (2026-07-08T06:43:42.605Z)

what should be the merge order between 527 and 528?

### 333. assistant (2026-07-08T06:44:22.578Z)

Let me verify the actual file overlap and conflict potential empirically rather than assume:

### 334. user (2026-07-08T06:44:27.465Z)

```
=== #527 ===
{"base":"master","files":["backend/pipeline/stitcher/compiler.py","tests/test_compiler.py"],"head":"fix/nvenc-probe-resolution","mergeable":"MERGEABLE","state":"OPEN"}
=== #528 ===
{"base":"master","files":["backend/orchestration.py","backend/worker/__main__.py","docs/COMPUTE_DECOUPLED_SERVING/PHASE_PLACEMENT.md","tests/test_compile_dispatch.py","tests/test_worker.py"],"head":"fix/521-worker-mirror-required","mergeable":"MERGEABLE","state":"OPEN"}
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 335. assistant (2026-07-08T06:44:54.966Z)

Both PRs are **MERGEABLE** and touch **completely disjoint files** (#527: `compiler.py` + `test_compiler.py``; #528: `orchestration.py`, `worker/__main__.py`, `PHASE_PLACEMENT.md`, `test_compile_dispatch.py`, `test_worker.py``). Let me prove no conflict empirically with a trial three-way merge:

### 336. user (2026-07-08T06:45:02.474Z)

```
* branch fix/nvenc-probe-resolution -> FETCH_HEAD
* branch fix/521-worker-mirror-required -> FETCH_HEAD
=== trial merge both branch trees onto master (exit 0 + no conflict markers = clean in ANY order) ===
merge-tree exit=0
0
=== overlapping files between the two PRs (should be NONE) ===
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 337. assistant (2026-07-08T06:45:17.111Z)

Empirically confirmed: **merge order doesn't matter.**

- `git merge-tree --write-tree` of the two branches ? exit 0, zero conflict markers.`
- **Zero overlapping files** ? #527 touches only `compiler.py` + `test_compiler.py``; #528 touches `orchestration.py`, `worker/__main__.py`, the doc, and two different test files.`
- **No semantic dependency:** #527 changes the NVENC *probe* resolution (which encoder gets selected); #528 changes *mirror-failure handling*. Neither reads the other's code ? and #528's tests stub `VideoCompiler`, so the probe is irrelevant to them.`

So: **merge them in whatever order is convenient, no rebase required.** Git merges per-file, and disjoint file sets can't textually conflict.

Two practical notes:

1. **After the first one merges**, GitHub recomputes the second's status ? it'll stay `MERGEABLE` and green (disjoint). Just glance at it before clicking; no rebase. A merge queue, if enabled, handles this automatically.`
2. **Both should land in the release.** They're independent but complementary on the cloud-serving compile path: #528 stops a silent 404-on-"success", #527 stops NVENC from silently no-op'ing to libx264 on the L4. Shipping #521 without #527 means the headline NVENC speedup may not actually engage.

If you want an arbitrary tiebreaker, take **#528 first** (higher-severity data-loss fix) then **#527** ? but it's genuinely a coin flip.

### 338. user (2026-07-08T06:46:53.978Z)

merged

### 339. assistant (2026-07-08T06:47:32.466Z)

Let me verify master is coherent post-merge and run the combined result through the suite ? since each PR's CI ran against a base \*without\* the other, a run against the actual merged commit closes that gap (the disjoint analysis predicts green, but a release deserves the real check).

### 340. user (2026-07-08T06:47:38.924Z)

```
=== master log ===
65ee4c6 fix(serving): probe NVENC at 256x256, not 64x64, so it isn't falsely rejected
(#521/#525) (#527)
008a5d4 fix(compile): fail the L4-worker stitch when the reel mirror-upload fails (#521 follow-
up) (#528)
dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521) (#525)
bf97700 docs(runlogs): record live SPA CUJ closing the 2 owner-gated validation-latency cells
($9) (#522)
699d85b feat(validation): cache validation verdict by video_id to skip re-validation on reupload
(#518)
=== reset my (merged) worktree to the released master ===
HEAD is now at 65ee4c6 fix(serving): probe NVENC at 256x256, not 64x64, so it isn't falsely
rejected (#521/#525) (#527)
=== both fixes present in master? ===
#528 require_mirror: backend/orchestration.py:6 backend/worker/__main__.py:1
#527 probe resolution (should NOT be 64x64):
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 341. assistant (2026-07-08T06:47:51.689Z)

Master has all three (#525 ? #528 ? #527). `require\_mirror` is present (6x in orchestration, 1x in worker). Let me confirm #527's probe change and start the full suite against the exact released commit:

### 342. user (2026-07-08T06:47:52.648Z)

```
=== #527's probe in merged master ===
35:#: 64x64 and 128x128 but accepts ~145px and up; 256x256 keeps a safe margin. Do NOT shrink
below
36:#: 256x256 ? guarded by tests/test_compiler.py::test_probe_encoder_uses_safe_resolution.
37:_PROBE_SOURCE_SIZE = "256x256"
40:def _probe_encoder_available(encoder: str) -> bool:
42: lavfi test encode to null at ``_PROBE_SOURCE_SIZE`` (a resolution ? NVENC's minimum, so
the
60: "lavfi",
62: f"color=c=black:s={_PROBE_SOURCE_SIZE}:d=0.1",
276: box = f":box=1:boxcolor=black@{min(opacity, 0.5):.2f}:boxborderw=10"
279: f":fontcolor=white@{opacity:.2f}:fontsize=h/{divisor}:{dt_xy}{box}"
Shell cwd was reset to C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer
```

### 343. user (2026-07-08T06:47:52.694Z)

Command running in background with ID: bihc5yblz. Output is being written to:  
C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\75d3d121-f240-4293-afa5-b436a379ad31\tasks\bihc5yblz.output. You will be notified when it completes. To check interim output, use Read on that file path.

### 344. assistant (2026-07-08T06:48:16.591Z)

#527's fix is sound: `\_PROBE\_SOURCE\_SIZE = "256x256"` with a guard test and a clear comment (NVENC rejects 64/128px, accepts ~145px+). The full suite is running against the exact released commit `65ee4c6` in the background.

Let me update the memory index to reflect the merges (re-reading first, since the sibling

session may have touched it):

#### 345. user (2026-07-08T06:48:17.464Z)

```
1 # Memory index
2
3 - [PR #525 ? compile-on-L4 (#521)](pr-525-compile-on-l4.md) ? config-driven
phase?machine placement + L4 NVENC stitch; Part 3 (cheaper watermark) deferred.
4 - [PR #527 ? NVENC probe 64x64?256x256](pr-527-nvenc-probe-resolution.md) ? sub-minimum
probe silently defeated #521's NVENC-on-L4 win; MERGED to master (65ee4c6); STILL awaits owner
live-GPU confirm; monitoring follow-up #529.
5 - [PR #528 ? worker mirror-required](pr-528-worker-mirror-required.md) ? #521 follow-up:
L4-worker stitch now fails on a lost reel mirror-upload (was silent 404-on-success); found by
adversarial review, CI green.
6
```

#### 346. user (2026-07-08T06:48:28.277Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 347. user (2026-07-08T06:48:33.265Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\memory\pr-528-worker-mirror-required.md has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 348. assistant (2026-07-08T06:48:47.485Z)

Master is coherent ? `65ee4c6` carries all three (#525 ? #528 ? #527), both follow-up fixes are present, and #527's probe fix is sound (256x256 with a guard test). Memory updated.

I'm running the full suite against that exact released commit now (in the background) to close the one gap CI didn't cover ? each PR's CI ran against a base \*without\* the other, so this is the first run with all three combined. The disjoint-file analysis predicts green; I'll confirm when it finishes.

**\*\*Release baseline summary:\*\***

```
| commit | PR | what |
|---|---|---|
| `dac6736` | #525 | #521 ? config-driven placement + compile-on-L4 + NVENC |
| `008a5d4` | #528 | fail the L4 stitch on a lost reel mirror (was silent 404-on-success) |
| `65ee4c6` | #527 | probe NVENC at 256x256 so it isn't falsely rejected on the L4 |
```

**\*\*One thing to keep in the release notes:\*\*** the real **\*\*L4 + NVENC** execution is still **unverified** ? the deployed `rally-worker:latest` predates the `compile` subcommand, so the compile-on-L4 path first runs for real only once you rebuild the worker image from this master. I validated the full worker path against real GCS locally (fetch ? stitch ? upload reel + marker), and both follow-up bugs were found/fixed pre-deploy, but the NVENC hardware encode itself will get its first real exercise on that rebuild. Worth a smoke check on the first cloud compile after deploy.

I'll report back the moment the suite lands.

#### 349. user (2026-07-08T06:51:42.466Z)

```
<task-notification>
<task-id>bihc5yblz</task-id>
<tool-use-id>toolu_01WqYDTXQdCdNfT55Vzj5fbN</tool-use-id>
<output-file>C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\75d3d121-f240-4293-afa5-b436a379ad31\tasks\bihc5yblz.output</output-file>
```

```
<status>completed</status>
<summary>Background command "Run full suite against merged master" completed (exit code
0)</summary>
</task-notification>
```

### 350. assistant (2026-07-08T06:51:46.975Z)

The suite finished with **exit 0**. Let me confirm the numbers:

### 351. user (2026-07-08T06:51:49.133Z)

```
SUITE_EXIT=0
Required test coverage of 70% reached. Total coverage: 85.82%
1882 passed, 8 skipped, 25 warnings in 225.47s (0:03:45)
=====
RESULT: PASS - all tests green
=====
```

### 352. assistant (2026-07-08T06:51:58.660Z)

Full suite **green** against the exact released commit ``65ee4c6``: **1882 passed, 8 skipped, coverage 85.82%** ? the first run with #525 + #528 + #527 all combined (which no individual PR's CI had verified). The release baseline is confirmed coherent.

Let me clean up my two now-merged worktrees:

### 353. user (2026-07-08T06:52:02.914Z)

```
error: failed to delete '.git/worktrees/kg-wt525fix': Permission denied
removed C:/kg-wt525fix
error: failed to delete '.git/worktrees/kg-wt521': Permission denied
removed C:/kg-wt521
=== my worktrees now ===
C:/Users/avidu/Projects/khelsutra-guru/bhi-wt-cheap-watermark
dac6736 [perf/521-cheap-watermark]
```

### 354. assistant (2026-07-08T06:52:33.553Z)

Everything's wrapped up. Cleanup note: the worktree `*directories*` were removed, but Git couldn't delete their ``.git/worktrees/*`` admin entries (``Permission denied`` ? the shared ``.git`` is locked by concurrent sessions, same contention we saw earlier). Those dangling entries are harmless and a ``git worktree prune`` will clear them once the other sessions release their locks. I left the sibling ``bhi-wt-cheap-watermark`` worktree (the watermark task you started) untouched.

## Release baseline ? verified and coherent ?

Master ``65ee4c6`` carries all three, and the **full suite passes** against that exact commit<sup>2</sup> (1882 passed, 8 skipped, coverage 85.82%) ? the first run exercising #525 + #528 + #527 together, which no individual PR's CI had done:

```
| commit | PR | what |
|---|---|---|
| `dac6736` | #525 | #521 ? config-driven phase placement + compile-on-L4 + NVENC |
| `008a5d4` | #528 | fail the L4 stitch on a lost reel mirror (was a silent 404-on-"success") |
| `65ee4c6` | #527 | probe NVENC at 256x256 so it isn't falsely rejected on the L4 |
```

## For your release notes / first post-deploy check

- **Real L4 + NVENC** is still unverified in prod<sup>2</sup> ? the deployed ``rally-worker:latest`` predated the ``compile`` subcommand, so the `compile-on-L4` path runs for real only after you rebuild the worker image from this master (``gcloud builds submit --config=deploy/cloudrun/cloudbuild.yaml``). I validated the whole worker path against real GCS locally; the hardware NVENC encode gets its first real exercise on that rebuild.
- **Smoke check after deploy:** drive one cloud compile and confirm (a) the reel lands at

`gs://?/highlights/<video\_id>/?` and (b) the worker log shows `encoding clips with h264\_nvenc` (not the libx264 fallback) ? that confirms both #527 (probe passes) and the #521 speedup actually engaged.

- Still open and independent: the **\*\*cheaper-watermark\*\*** follow-up (your `perf/521-cheap-watermark` session) ? the deeper win where the stitch avoids a full re-encode entirely.

Net: #521 shipped with two review-caught bugs fixed before they could bite, and the release baseline is green.