

claude-windows / 75d3d121-f240-4293-afa5-b436a379ad31

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\75d3d121-f240-4293-afa5-b436a379ad31\subagents\workflows\wf_5f13fc79-bad\agent-ab5995df7f12953f4.jsonl`  
- SHA-256: `8d1cfd8fe680531f5131ba9f17b15f24b353fce39f51d8fc7f4e1ad40e3bea11`  
- Source modified: `2026-07-08T06:18:05+00:00`  
- Imported at: `2026-07-08T15:59:28+00:00`  
- project: `wf_5f13fc79-bad`  
- session_id: `75d3d121-f240-4293-afa5-b436a379ad31`
```

Transcript

1. user (2026-07-08T06:14:14.437Z)

CONTEXT: This is GitHub issue #521 (Khelsutra/badminton-highlight-indexer), MERGED as PR #525. It makes pipeline phase?machine placement config-driven and dispatches the reel compile/stitch to an L4 Cloud Run worker (NVENC), mirroring the existing detection dispatch.

The full unified diff is at: undefined (read it first).

The full post-change source is under: undefined/ (read whole files there for context ? e.g. undefined/backend/orchestration.py, undefined/backend/worker/__main__.py, undefined/backend/compute/cloud_run.py, undefined/backend/config/placement.py, undefined/backend/config/models.py, undefined/backend/pipeline/stitcher/compiler.py).

Key design invariants that MUST hold (a violation is a real bug):

- DEFAULT-OFF: with no phase_placement config and compute_target unset, behaviour is byte-identical to before #521 (validation+compile on control-plane, detect in-process).
- The cloud compile path must return the SAME {status, filepath, download_url, video_id} shape the SPA already polls for; a stitch failure must be a clean failed-job, never a crash.
- The worker 'compile' subcommand reads a bucket-staged spec (resolved time-pairs + stitching config + source uri) and needs NO control-plane DB. It reuses orchestration.__stitch_reel (shared core).
- Encoder: all clips in one reel must share ONE codec so the concat stays -c copy. NVENC is probed once and falls back to libx264 if unavailable (CI/non-GPU) ? the fallback must be all-or-nothing, never per-clip (which would break the concat).
- The config guard rejects validation=cloud_run_l4 without a cloud segment (nobody-validates footgun).
- Marker prefixes: compile uses _compile/, detect uses _dispatch/ ? must not collide.

Only report DEFECTS in the changed code (logic bugs, broken invariants, edge cases that produce a wrong result / crash / silent data loss, or missing test coverage of a real failure mode). Do NOT report style, naming, or speculative concerns. Cite exact file:line from undefined.

A reviewer claims this DEFECT in the merged #521 code. Try HARD to REFUTE it by reading the actual code at undefined. Default to is_real=false if you cannot concretely confirm the failure. Skeptic #2.

CLAIM (high/test-coverage) backend/orchestration.py:1263

summary: No test exercises the L4 worker's real `_stitch_reel` mirror-upload; a swallowed mirror failure there reports success (rc 0) but the reel is lost, so the SPA gets a 404 for a job it was told succeeded.

failure scenario: On the L4, a compile job runs `cmd_compile -> orchestration._stitch_reel`. The `ffmpeg` stitch succeeds locally, then ``storage.put(out_path, gs://.../highlights/<vid>/<name>`` fails (transient GCS 5xx, throttling, or a creds/permission hiccup). The ``except Exception`` at `orchestration.py:1263` swallows it (logs a warning) and `_stitch_reel` returns ``status="success"`` with the intended ``reel_uri`` (line 1270-1276). `cmd_compile` (`worker/__main__.py:513-514`) sets `rc=0` purely from that status and writes a SUCCESS done-marker; `_dispatch_compile_remote` returns `{status: success, download_url}`. But the reel was never uploaded and the worker's local file is discarded when the Job exits, so `/api/highlights -> signed_reel_url -> storage.exists()` is False, the control-plane has no local fallback (the stitch ran on the worker), and the download 404s forever. Every #521 test hides this: `test_worker.py`'s `cmd_compile` tests monkeypatch `orchestration._stitch_reel`, and `test_compile_dispatch.py` / `test_compute_backend.py` use a fake `CloudRunJobClient` that writes the marker directly, so the real worker mirror path (and its swallow-on-failure) is never run.

2. assistant (2026-07-08T06:14:24.891Z)

I'll investigate this claim by reading the actual code. Let me start with the key files.

3. user (2026-07-08T06:14:26.256Z)

```
1180 def signed_reel_url(cfg, video_id: str, name: str) -> Optional[str]:
1181     """A short-lived signed URL for a reel mirrored to the output bucket, so
/api/highlights can
1182     307-redirect off the app's response path (bypassing Cloud Run's ~32 MiB cap + the
ephemeral local
1183     /tmp; #463). None ? no bucket copy ? serve the local file. Uses the active storage
config for BOTH
1184     the existence check and the signing (extracted here to keep backend/main.py lean ?
P23/P24)."""
1185     reel_uri = reel_object_uri(cfg, video_id, name)
1186     if not reel_uri:
1187         return None
1188     settings = _storage_settings(cfg)
1189     if storage.exists(reel_uri, config=settings):
1190         return storage.signed_url(reel_uri, config=settings, method="GET")
1191     return None
1192
1193
1194 def _compile_wait_message(elapsed_sec: float) -> str:
1195     """The live job.progress line while an L4 compile/stitch is in flight (#521; mirrors
1196     ``_wait_progress_message``). Minute-granular; under a minute reads as cold-start."""
1197     minutes = int(elapsed_sec) // 60
1198     if minutes < 1:
1199         return "stitching on cloud GPU (NVENC) ? starting up"
1200     return f"stitching on cloud GPU (NVENC) ? {minutes} min elapsed"
1201
1202
1203 def _stitch_reel(
1204     cfg,
1205     stitching,
1206     video_id: str,
1207     source_ref: str,
1208     time_pairs: List[Tuple[float, float]],
1209     out_name: str,
1210     locale: Optional[str],
1211     notify,
1212 ) -> Dict[str, Any]:
1213     """Stitch + mirror-upload core, shared by the control-plane compile handler AND the
L4 worker
1214     (#521). ``time_pairs`` are already RESOLVED (the caller filtered via
``_resolve_compile``), so this
```

```

1215         needs NO control-plane DB ? the worker reads them from the staged spec. Builds the
VideoCompiler
1216         with ``stitching.encoder`` (NVENC on the L4, libx264 on the CP), fetches the source,
runs the
1217         trim/watermark/concat, mirrors the reel to
``<output_root>/highlights/<video_id>/<name>``, and
1218         returns the ``{status, filepath, download_url, video_id, reel_uri}`` the SPA polls
for. Returns a
1219         ``status='failed'`` result on a stitch error (never raises for an expected ffmpeg
failure)."""
1220         import backend.main as _main # call-time seam: VideoCompiler stays patchable on
backend.main
1221
1222         <<<<<<< Updated upstream
1223         =====
1224         notify = progress or (lambda stage: None)
1225         params = json.loads(job.get("params") or "{}")
1226         video_id = job.get("video_id") or ""
1227         segment_ids = params.get("segment_ids") or []
1228         locale = params.get("locale")
1229         rendition = str(params.get("rendition") or "original")
1230         if rendition not in {"original", "sd"}:
1231             return {
1232                 "status": "failed",
1233                 "message": f"Unsupported compile rendition: {rendition}",
1234                 "video_id": video_id,
1235             }
1236
1237         notify("resolving")
1238         try:
1239             video, kept, out_name = _resolve_compile(
1240                 db, cfg, video_id, segment_ids, params.get("output_filename", "")
1241             )
1242         except _CompileError as e:
1243             return {"status": "failed", "message": e.detail, "video_id": video_id}
1244
1245         stitching = getattr(cfg, "stitching", None) or StitchingConfig()
1246         >>>>>>> Stashed changes
1247         output_dir = (
1248             getattr(getattr(cfg, "output", None), "output_dir", "output") or "output"
1249         )
1250         # Localize the branding watermark to the UI locale the SPA rendered in (recorded on
the job);
1251         # unknown/absent locale keeps the configured default unchanged.
1252         localized_watermark = localize_watermark(stitching.watermark, locale)
1253         encode_profile = _rendition_encode_profile(stitching, rendition)
1254         compiler = _main.VideoCompiler(
1255             output_dir,
1256             begin_padding=stitching.begin_padding,
1257             end_padding=stitching.end_padding,
1258             watermark=localized_watermark,
1259         <<<<<<< Updated upstream
1260         encoder=getattr(stitching, "encoder", "libx264"),
1261         =====
1262         encode_profile=encode_profile,
1263         >>>>>>> Stashed changes
1264         )
1265
1266         notify("stitching")
1267         logger.info(
1268             "[compile] stitching video_id=%s rendition=%s output=%s",
1269             video_id,
1270             rendition,
1271             out_name,
1272         )

```

```

1273     local_src = ""
1274     try:
1275         # The stored `filepath` is the SOURCE ref (gs://, s3://, gdrive://, or a local
1276         # path). Resolve to
1277         # a LOCAL file the stitcher can read ? identity for a local path, fetch-to-
1278         # scratch for a URI (the
1279         # same seam the process path uses); without it, compiling a bucket-ingested
1280         # video can't open gs://.
1281         local_src = storage.acquire_local_input(
1282             source_ref, getattr(cfg, "storage", None)
1283         )
1284         # The ffmpeg stitch is CPU-heavy ? serialize it so a raised drain worker count
1285         # (for remote
1286         # dispatch) doesn't run multiple compiles at once on the box (see
1287         # _LOCAL_COMPUTE_LANE).
1288         with _LOCAL_COMPUTE_LANE:
1289             out_path = compiler.compile_highlights(local_src, time_pairs, out_name)
1290         except Exception as e: # noqa: BLE001 ? a stitch/ffmpeg failure is a clean job
1291             # FAILURE, not a crash
1292             logger.error(
1293                 "[compile] stitching failed for video_id=%s: %s", video_id, e, exc_info=True
1294             )
1295             return {
1296                 "status": "failed",
1297                 "message": f"Highlight stitching failed: {e}",
1298                 "video_id": video_id,
1299             }
1300         finally:
1301             if local_src:
1302                 storage.cleanup_acquired_local_input(
1303                     source_ref, local_src, getattr(cfg, "storage", None)
1304                 )
1305             # #463: mirror the reel to the output bucket so /api/highlights can 307 to a signed
1306             # URL ? Cloud Run
1307             # caps a buffered response at ~32 MiB (a real reel is bigger) and the local /tmp
1308             # copy dies when the
1309             # instance recycles. A mirror-upload failure must NOT fail an otherwise-good stitch
1310             # (the endpoint
1311             # still falls back to the local file on the same instance). On the L4 worker the
1312             # mirror is how the
1313             # reel reaches the bucket at all (the worker's local file is discarded when the Job
1314             # exits).
1315             reel_uri = reel_object_uri(cfg, video_id, out_name)
1316             if reel_uri:
1317                 try:
1318                     notify("uploading reel")
1319                     # Thread the active storage settings so the configured backend (S3
1320                     # endpoint/region, GCS
1321                     # project/client) is used, not a fresh default.
1322                     storage.put(out_path, reel_uri, config=_storage_settings(cfg))
1323                 except Exception as e: # noqa: BLE001 ? best-effort mirror; local stitch
1324                     # already succeeded
1325                     logger.warning("[compile] could not mirror reel to %s: %s", reel_uri, e)
1326             # `filepath` is server-local (not browser-reachable); also return the download URL
1327             # so the SPA needn't
1328             # hand-build it (P6 / PER_RALLY_SELECTION.md §4.3).
1329             download_url = f"/api/highlights/{quote(video_id)}/{quote(out_name)}"
1330             notify("finished")
1331             return {

```

4. user (2026-07-08T06:14:26.612Z)

backend\worker__main__.py

5. user (2026-07-08T06:14:27.058Z)

tests\test_worker.py

6. user (2026-07-08T06:14:27.259Z)

tests\test_compile_dispatch.py

7. assistant (2026-07-08T06:14:56.206Z)

I notice the working copy of `orchestration.py` has git merge-conflict markers (a stash conflict). I need the clean merged version. Let me check git state and read the committed version.

8. user (2026-07-08T06:14:58.986Z)

On branch codex/issue-523-sd-reel-plan
Your branch is up to date with 'origin/master'.

Changes to be committed:

```
(use "git restore --staged <file>..." to unstage)
    modified:   README.md
    modified:   SESSION_HANDOFF.md
    modified:   backend/config/models.py
    modified:   backend/infrastructure/database_jobs.py
    modified:   backend/main.py
    modified:   config.json
    modified:   tests/test_api_endpoints.py
    modified:   tests/test_async_jobs.py
    modified:   tests/test_compiler.py
    modified:   tests/test_watermark.py
```

Unmerged paths:

```
(use "git restore --staged <file>..." to unstage)
(use "git add <file>..." to mark resolution)
    both modified:   backend/orchestration.py
    both modified:   backend/pipeline/stitcher/compiler.py
```

Untracked files:

```
(use "git add <file>..." to include in what will be committed)
    score_golden_gate.py
    score_local.py
```

---LOG---

```
dac6736 perf(serving): run compile/stitch on the L4 worker; config-driven phase?machine
placement (#521) (#525)
bf97700 docs(runlogs): record live SPA CUJ closing the 2 owner-gated validation-latency cells
($9) (#522)
699d85b feat(validation): cache validation verdict by video_id to skip re-validation on reupload
(#518)
71c88ef docs(runlogs): validation-latency closure report for #512 + #513 (cloud-serving) (#519)
115b83d docs(runlogs): record live smoke of the crash-before-marker fast-fail (#515) (#517)
ef41473 fix(cloud-run): fail fast when a worker execution terminates without a done-marker
(#515)
f5c9dd9 chore(ui): re-vendor SPA bundle with the hosted compute-picker cloud default (#516)
ec2d080 fix(serving): fail fast when compute=local forces weightless in-process detection (#511)
---STASH---
stash@{0}: On codex/issue-523-sd-reel-plan: sd-reel-engine-wip
```

9. assistant (2026-07-08T06:15:05.473Z)

The working tree has a stash-pop conflict on top of `dac6736` (PR #525 = merged #521). Let me read the clean committed versions from HEAD.

10. user (2026-07-08T06:15:08.187Z)

```
def _compile_wait_message(elapsed_sec: float) -> str:
    """The live job.progress line while an L4 compile/stitch is in flight (#521; mirrors
    ``_wait_progress_message``). Minute-granular; under a minute reads as cold-start."""
    minutes = int(elapsed_sec) // 60
    if minutes < 1:
        return "stitching on cloud GPU (NVENC) ? starting up"
    return f"stitching on cloud GPU (NVENC) ? {minutes} min elapsed"

def _stitch_reel(
    cfg,
    stitching,
    video_id: str,
    source_ref: str,
    time_pairs: List[Tuple[float, float]],
    out_name: str,
    locale: Optional[str],
    notify,
) -> Dict[str, Any]:
    """Stitch + mirror-upload core, shared by the control-plane compile handler AND the L4
    worker
    (#521). ``time_pairs`` are already RESOLVED (the caller filtered via ``_resolve_compile``),
    so this
    needs NO control-plane DB ? the worker reads them from the staged spec. Builds the
    VideoCompiler
    with ``stitching.encoder`` (NVENC on the L4, libx264 on the CP), fetches the source, runs
    the
    trim/watermark/concat, mirrors the reel to ``<output_root>/highlights/<video_id>/<name>``,
    and
    returns the ``{status, filepath, download_url, video_id, reel_uri}`` the SPA polls for.
    Returns a
    ``status='failed'`` result on a stitch error (never raises for an expected ffmpeg
    failure)."""
    import backend.main as _main # call-time seam: VideoCompiler stays patchable on
    backend.main

    output_dir = (
        getattr(getattr(cfg, "output", None), "output_dir", "output") or "output"
    )
    # Localize the branding watermark to the UI locale the SPA rendered in (recorded on the
    job);
    # unknown/absent locale keeps the configured default unchanged.
    localized_watermark = localize_watermark(stitching.watermark, locale)
    compiler = _main.VideoCompiler(
        output_dir,
        begin_padding=stitching.begin_padding,
        end_padding=stitching.end_padding,
        watermark=localized_watermark,
        encoder=getattr(stitching, "encoder", "libx264"),
    )

    notify("stitching")
    local_src = ""
    try:
        # The stored `filepath` is the SOURCE ref (gs://, s3://, gdrive://, or a local path).
        Resolve to
        # a LOCAL file the stitcher can read ? identity for a local path, fetch-to-scratch for a
        URI (the
        # same seam the process path uses); without it, compiling a bucket-ingested video can't
        open gs://.
        local_src = storage.acquire_local_input(
            source_ref, getattr(cfg, "storage", None)
        )
    )
```

```

        # The ffmpeg stitch is CPU-heavy ? serialize it so a raised drain worker count (for
remote
        # dispatch) doesn't run multiple compiles at once on the box (see _LOCAL_COMPUTE_LANE).
        with _LOCAL_COMPUTE_LANE:
            out_path = compiler.compile_highlights(local_src, time_pairs, out_name)
        except Exception as e: # noqa: BLE001 ? a stitch/ffmpeg failure is a clean job FAILURE, not
a crash
            logger.error(
                "[compile] stitching failed for video_id=%s: %s", video_id, e, exc_info=True
            )
            return {
                "status": "failed",
                "message": f"Highlight stitching failed: {e}",
                "video_id": video_id,
            }
        finally:
            if local_src:
                storage.cleanup_acquired_local_input(
                    source_ref, local_src, getattr(cfg, "storage", None)
                )

        # #463: mirror the reel to the output bucket so /api/highlights can 307 to a signed URL ?
Cloud Run
        # caps a buffered response at ~32 MiB (a real reel is bigger) and the local /tmp copy dies
when the
        # instance recycles. A mirror-upload failure must NOT fail an otherwise-good stitch (the
endpoint
        # still falls back to the local file on the same instance). On the L4 worker the mirror is
how the
        # reel reaches the bucket at all (the worker's local file is discarded when the Job exits).
        reel_uri = reel_object_uri(cfg, video_id, out_name)
        if reel_uri:
            try:
                notify("uploading reel")
                # Thread the active storage settings so the configured backend (S3 endpoint/region,
GCS
                # project/client) is used, not a fresh default.
                storage.put(out_path, reel_uri, config=storage_settings(cfg))
            except Exception as e: # noqa: BLE001 ? best-effort mirror; local stitch already
succeeded
                logger.warning("[compile] could not mirror reel to %s: %s", reel_uri, e)

        # `filepath` is server-local (not browser-reachable); also return the download URL so the
SPA needn't
        # hand-build it (P6 / PER_RALLY_SELECTION.md §4.3).
        download_url = f"/api/highlights/{quote(video_id)}/{quote(out_name)}"
        notify("finished")
        return {
            "status": "success",
            "filepath": out_path,
            "download_url": download_url,
            "video_id": video_id,
            "reel_uri": reel_uri or "",
        }

def _dispatch_compile_remote(
    cfg,
    db,
    video_id: str,
    segment_ids: List[int],
    output_filename: str,
    locale: Optional[str],
    notify,
) -> Dict[str, Any]:

```

```

    """#521: run the compile/stitch on the L4 worker instead of the CPU-only control-plane. The
    control-plane owns the DB, so it RESOLVES the reel's segments here (`_resolve_compile`),
    stages a
    small compile-spec (resolved time-pairs + the resolved stitching/watermark config + the
    source URI)
    to the bucket, dispatches a `compile` Cloud Run Job on the L4 (NVENC), bucket-polls the
    done-marker, and returns the same `{status, filepath, download_url}` shape the SPA
    expects. The
    worker needs NO control-plane DB ? everything it needs rides the spec. Mirrors
    `_maybe_dispatch_remote` (the DETECT dispatch)."""
    from backend.compute import CompileJobSpec, get_compute_backend

    notify("resolving")
    try:
        video, kept, out_name = _resolve_compile(
            db, cfg, video_id, segment_ids, output_filename

```

11. assistant (2026-07-08T06:15:29.934Z)

This is the merged code. The swallow-and-return-success logic is confirmed. Now the crux: does the worker `cmd\_compile` or the control-plane `\_dispatch\_compile\_remote` verify the reel actually landed in the bucket before declaring success? And do the tests really mock out the real mirror path? Let me read the worker and dispatch code, plus the tests.

12. user (2026-07-08T06:15:31.562Z)

```

1      """Worker dispatcher ? pull from storage ? run the pipeline ? push outputs back.
2
3      python -m backend.worker infer --video-uri <uri> --out-uri <prefix> \
4          [--segmenter wasb_hybrid] [--config config.json] [--set k=v ...] [--scratch DIR]
5      python -m backend.worker train ...          # Gen-0 train-from-bucket (follow-up)
6
7      The URI scheme selects the storage backend via `backend.storage` (a bare path /
`local://`
8      is an identity passthrough; `s3://` / `gcs://` is an object store), so existing
local runs
9      are byte-for-byte unchanged. Combine with `RALLY_DETECTOR_IMPL=stub` and
10     `--set indexing.skip_ai_handoff=true` for a **secret-free, GPU-free** run (no Gemini
key) ?
11     the path validated end-to-end on DigitalOcean (`docs/SETUP_NEW_MACHINE.md`).
12
13     This is a thin orchestrator: it composes the SAME building blocks the `backend.main`
CLI uses
14     (validator registry ? segmenter ? report), never a forked pipeline.
15     """
16
17     from __future__ import annotations
18
19     import argparse
20     import csv
21     import json
22     import logging
23     import os
24     import re
25     import sys
26     import tempfile
27     from typing import Any, List
28
29     from backend import storage
30     from backend.compute import ComputeJobSpec, get_compute_backend
31     from backend.config import (
32         StartupCheckError,
33         enforce_startup_checks,
34         load_config,
35         probe_cuda_available,

```



```

36 )
37 from backend.infrastructure.database import Database
38 from backend.pipeline.segmenters.base import segmenter_registry
39 from backend.pipeline.validators.base import registry as validator_registry
40 from backend.reporting import emit_indexer_report
41 from backend.results import Failure
42 from backend.utils.ffmpeg import ffprobe_bin # P2a: single ffmpeg/ffprobe resolver
43 from backend.utils.hashing import compute_video_id
44
45 logger = logging.getLogger("backend.worker")
46
47
48 def _coerce(v: str) -> Any:
49     """Coerce a ``--set`` string value to bool/int/float where it round-trips, else
str."""
50     low = v.lower()
51     if low in ("true", "false"):
52         return low == "true"
53     for cast in (int, float):
54         try:
55             return cast(v)
56         except ValueError:
57             pass
58     return v
59
60
61 def _apply_overrides(config: Any, sets: List[str]) -> None:
62     """Apply ``a.b.c=value`` overrides onto the typed config (e.g.
indexing.skip_ai_handoff=true)."""
63     for kv in sets or []:
64         key, sep, val = kv.partition("=")
65         if not sep:
66             raise ValueError(f"--set expects k=v, got {kv!r}")
67         parts = key.split(".")
68         obj = config
69         for p in parts[:-1]:
70             obj = getattr(obj, p)
71         setattr(obj, parts[-1], _coerce(val))
72
73
74 def _write_intervals_csv(segments: List[dict], path: str) -> None:
75     """Decimal-second ``[start,end)`` rallies + shot count / ending reason ? the join-
friendly
76     artifact (same schema as the rally-annotator labels)."""
77     with open(path, "w", newline="") as f:
78         w = csv.writer(f)
79         w.writerow(["start", "end", "shot_count", "ending_reason"])
80         for s in segments:
81             w.writerow(
82                 [
83                     f"{float(s.get('start_time', 0.0)):.3f}",
84                     f"{float(s.get('end_time', 0.0)):.3f}",
85                     s.get("shot_count", ""),
86                     s.get("ending_reason", s.get("shot_count_confidence", "")),
87                 ]
88             )
89
90
91 def _write_report_json(
92     video_id: str, segments: List[dict], status: str, path: str, video_uri: str = ""
93 ) -> None:
94     with open(path, "w") as f:
95         json.dump(
96             {
97                 "video_id": video_id,

```

```

98         "status": status,
99         "segment_count": len(segments),
100         # The SOURCE the detection ran on (CP-rebuild P3, #485): a restart-wiped
101         # control plane re-hydrating from outputs/<md5>/ restores the video row
102         # with its truthful path ? without this, only the segments recover.
103         "video_uri": video_uri,
104         "segments": [
105             {
106                 "start": float(s.get("start_time", 0.0)),
107                 "end": float(s.get("end_time", 0.0)),
108             }
109             for s in segments
110         ],
111     },
112     f,
113     indent=2,
114 )
115
116
117 def _serialize_and_push_outputs(
118     scratch: str,
119     out_uri: str,
120     video_id: str,
121     segments: List[dict],
122     status: str,
123     storage_cfg: dict,
124     video_uri: str = "",
125 ) -> List[str]:
126     """Write intervals.csv + report.json under ``<scratch>/outputs/<video_id>/`` and
push both to
127     ``<out_uri>/outputs/<video_id>/``; return the pushed URIs.
128
129     Shared by the SUCCESS path and the failure path (``_fail``) so a detector
timeout/abort leaves a
130     ``status="failed"`` report.json in the bucket ? the control plane's source of truth
131     (``orchestration.cached_process_result`` / ``probe_process_cache`` treat any
non-``success``
132     report as "no usable result" ? retry, and the restart-recovery poll waits for
report.json to
133     EXIST) ? instead of a missing artifact or a false 0-segment ``"success"``.
134     out_local = os.path.join(scratch, "outputs", video_id)
135     os.makedirs(out_local, exist_ok=True)
136     _write_intervals_csv(segments, os.path.join(out_local, "intervals.csv"))
137     _write_report_json(
138         video_id,
139         segments,
140         status,
141         os.path.join(out_local, "report.json"),
142         video_uri=video_uri,
143     )
144     out_prefix = out_uri.rstrip("/")
145     pushed: List[str] = []
146     for fn in sorted(os.listdir(out_local)):
147         uri = f"{out_prefix}/outputs/{video_id}/{fn}"
148         storage.put(os.path.join(out_local, fn), uri, config=storage_cfg)
149         pushed.append(uri)
150     return pushed
151
152
153 def _run_startup_checks(config: Any) -> bool:
154     """M5 per-profile startup checks for the worker (the compute process, so it probes
CUDA
155     best-effort). Logs warnings; returns False on a FATAL coherence error so the caller
aborts
156     cleanly (rc 2) rather than run an expensive job under an incoherent profile. A true

```

```

no-op
157         (returns True, ZERO work) when no deployment.profile is selected.
158
159         The CUDA probe is gated behind an active profile: ``probe_cuda_available()`` lazily
imports
160         torch, and pre-M5 the worker's infer/train path was torch-free (detection is
delegated to a
161         WSL/native subprocess). Probing only when a profile is set keeps the default-OFF
path a true
162         no-op of work, not just of output.
163         profile = getattr(getattr(config, "deployment", None), "profile", "") or ""
164         cuda = (
165             probe_cuda_available() if profile else None
166         ) # torch-free on the default-OFF path
167         try:
168             enforce_startup_checks(config, cuda_available=cuda, logger=logger)
169             return True
170         except StartupCheckError:
171             return False # already logged at ERROR by enforce_startup_checks
172
173
174     def _dispatch_remote(compute: Any, args: argparse.Namespace) -> int:
175         """M6: hand ONE infer job to a remote ComputeBackend (Cloud Run) and return its rc
once the
176         outputs are durably in the bucket (the backend bucket-polls the done-marker). The
dispatched
177         container runs this SAME cmd_infer INLINE (compute_target forced local_gpu) ? never
re-dispatch.
178
179         A remote job that never writes a completion marker exits via one of two CLEAN codes
(distinct
180         from the local 2/3), never a raw traceback:
181         * rc 4 ? the bucket poll exhausted its budget (``PolicyTimeout``): the execution
may still be
182         running (the shim best-effort cancels it).
183         * rc 5 ? the execution has TERMINATED without a marker
(``RemoteExecutionFailed``): a worker
184         crash/OOM/pre-marker abort caught FAST by the Executions-API probe, so instead
of the
185         ~65-min opaque wait the operator gets an immediate, log-pointing failure.
186         from backend.compute import RemoteExecutionFailed
187         from backend.infrastructure.execution_policy import PolicyTimeout
188
189         spec = ComputeJobSpec(
190             video_uri=args.video_uri,
191             out_uri=args.out_uri,
192             segmenter=args.segmenter,
193             config_overrides=tuple(args.set or ()),
194             skip_validation=bool(getattr(args, "skip_validation", False)),
195         )
196         try:
197             result = compute.run_infer(spec)
198         except RemoteExecutionFailed as e:
199             logger.error(
200                 "[worker] remote execution terminated without a completion marker: %s", e
201             )
202             return 5
203         except PolicyTimeout as e:
204             logger.warning(
205                 "[worker] remote compute did not complete (no marker within budget): %s", e
206             )
207             return 4
208         logger.info(
209             "[worker] remote compute done: video_id=%s rc=%d outputs=%s",
210             result.video_id,

```

```

211         result.returncode,
212         result.outputs_uri,
213     )
214     return result.returncode
215
216
217     def _write_done_marker(
218         done_uri: str,
219         video_id: str,
220         returncode: int,
221         segment_count: int,
222         status: str,
223         storage_cfg: dict,
224         scratch: str,
225     ) -> None:
226         """M6: write a tiny JSON completion marker to ``done_uri`` (via the storage layer)
for a remote
227         dispatcher's bucket-poll. Best-effort ? never raises into the worker's success
return."""
228         try:
229             marker = {
230                 "video_id": video_id,
231                 "returncode": returncode,
232                 "segment_count": segment_count,
233                 "status": status,
234             }
235             local = os.path.join(scratch, "_done.json")
236             with open(local, "w", encoding="utf-8") as f:
237                 json.dump(marker, f)
238             storage.put(local, done_uri, config=storage_cfg)
239             logger.info("[worker] wrote completion marker -> %s", done_uri)
240         except Exception as e: # never fail a good run because the marker push hiccuped
241             logger.warning("[worker] could not write done-marker %s: %s", done_uri, e)
242
243
244     def cmd_infer(args: argparse.Namespace) -> int:
245         config = load_config(args.config) if args.config else load_config()
246         _apply_overrides(config, args.set)
247         if not _run_startup_checks(config):
248             return 2
249
250         # M6: a cloud compute_target DISPATCHES to a Cloud Run Job instead of running in-
process.
251         # DEFAULT-OFF: get_compute_backend returns None for ""/local_gpu/cpu_mock ? fall
through to the
252         # unchanged inline path below. The dispatched container runs THIS cmd_infer inline
(the
253         # dispatcher forces compute_target=local_gpu), so it never recursively re-
dispatches.
254         compute = get_compute_backend(config)
255         if compute is not None:
256             return _dispatch_remote(compute, args)
257
258         storage_cfg = config.storage.model_dump()
259
260         scratch = args.scratch or tempfile.mkdtemp(prefix="rally-worker-")
261         os.makedirs(scratch, exist_ok=True)
262         config.output.output_dir = scratch
263
264         # 1. PULL ? local:// / bare path = identity passthrough; s3:// / gcs:// = download
to scratch.
265         local_video = storage.fetch(
266             args.video_uri, os.path.join(scratch, "in.mp4"), config=storage_cfg
267         )
268         print(f"[worker] pulled {args.video_uri} -> {local_video}")

```

```

269
270     # 2. IDENTITY ? whole-file MD5 (the labels?video join key; pulled bytes must be
byte-identical).
271     video_id = compute_video_id(local_video)
272
273     # 3. VALIDATE (same registry as the main CLI) ? UNLESS the dispatcher already did.
274     # On the cloud-serving path the CONTROL-PLANE is the authoritative validation gate:
it runs the
275     # SAME validators with ITS resolved config and only dispatches on PASS. Re-
validating here is
276     # redundant, decodes the whole video a SECOND time, and is outright WRONG when the
worker's config
277     # resolves different thresholds than the control-plane's ? observed 2026-07-07: a
5312x2988 clip
278     # (sharpness 11.9) PASSED the control-plane at validation.min_blur_threshold=8.0,
then the worker
279     # REJECTED it at its own default 20.0, so a SUCCESSFUL GPU dispatch failed after the
fact.
280     # --skip-validation (set by the dispatcher in orchestration._maybe_dispatch_remote)
skips it; the
281     # direct ``worker infer`` CLI has no control-plane, so it still validates by
default.
282     if getattr(args, "skip_validation", False):
283         print(
284             "[worker] validation SKIPPED (--skip-validation): the dispatcher is the
authoritative "
285             "gate and already validated this input."
286         )
287         val_results, val_context = [], {}
288     else:
289         val_results, val_context = validator_registry.run_all_with_context(
290             local_video, config
291         )
292     failures = [r for r in val_results if not r.passed]
293     db = Database(os.path.join(scratch, config.output.db_name))
294     db.add_video(
295         video_id,
296         local_video,
297         "failed" if failures else "processed",
298         [r.model_dump() for r in val_results],
299     )
300
301     def _fail(rc: int) -> int:
302         # A worker-side FAILURE (validation, unknown segmenter, or a detector
timeout/abort the
303         # segmenter surfaced as a Failure) must NOT masquerade as a 0-rally "success".
Write + push
304         # a status="failed" report.json/intervals.csv (0 segments) so the control
plane's bucket
305         # truth (submit-cache probe, restart-recovery) and the alpha-user result surface
see an
306         # explicit FAILURE ? the wf294 silent-failure fix. Best-effort: a push hiccup
must never
307         # mask the original failure. THEN the M6 done-marker (real rc + status) so a
remote
308         # dispatcher's bucket-poll terminates FAST instead of hanging until its budget
expires.
309         try:
310             _serialize_and_push_outputs(
311                 scratch,
312                 args.out_uri,
313                 video_id,
314                 [],
315                 "failed",
316                 storage_cfg,

```

```

317         str(getattr(args, "video_uri", "") or ""),
318     )
319 except Exception as e: # never mask the real failure on an artifact-push hiccup
320     logger.warning("[worker] could not push failure artifacts: %s", e)
321 # DEFAULT-OFF: no --done-uri ? no marker written (the unchanged local run).
322 if getattr(args, "done_uri", None):
323     _write_done_marker(
324         args.done_uri, video_id, rc, 0, "failed", storage_cfg, scratch
325     )
326 return rc
327
328 segments: List[dict] = []
329 segmenter_actual = None
330 segmentation_ran = False
331 status = "failed"
332 try:
333     if failures:
334         print(
335             "[worker] validation FAILED: "
336             + "; ".join(f"{r.validator_name}: {r.message}" for r in failures)
337         )
338         return _fail(2)
339 seg_name = args.segmenter or config.indexing.default_segmenter
340 segmenter_cls = segmenter_registry.get(seg_name)
341 if not segmenter_cls:
342     print(
343         f"[worker] unknown segmenter: {seg_name!r} "
344         f"(have {list(segmenter_registry.list_available())})"
345     )
346     return _fail(2)
347 segmenter_actual = seg_name
348 result = segmenter_cls(db, config).process_video(
349     video_id, local_video, max_frames=args.max_frames
350 )
351 segmentation_ran = True
352 if isinstance(result, Failure):
353     print(f"[worker] segmenter FAILED: {result.message}")
354     return _fail(3)
355 segments = result.value if hasattr(result, "value") else result
356 status = "success"
357 # Loud, never-silent diagnostic: a 0-segment success is almost always a
misconfig, not an
358 # empty match ? the cold-start failure mode (the substrate detector yielded no
usable
359 # trajectories). Explain it so a run is analysable from the log alone (re-run -v
for more).
360 if not segments:
361     detector_impl = (
362         os.environ.get("RALLY_DETECTOR_IMPL")
363         or getattr(config.indexing, "detector_impl", None)
364         or "auto"
365     )
366     logger.warning(
367         "infer produced 0 segments for video_id=%s (segmenter=%s,
detector_impl=%s). "
368         "The detector substrate yielded no usable trajectories/windows ? check
the detector "
369         "logs above; common causes: a missing ${AI_PROVIDER}_API_KEY with "
370         "skip_ai_handoff=False (the segmenter bails BEFORE detection ? set the
key or "
371         "indexing.skip_ai_handoff=true; cloud-serving now fails this at
startup), native "
372         "runner UNAVAILABLE (weights/repo/WASB_PYTHON env), the WASB env
produced no "
373         "detections, or the input resolution differs from the tuned proxy

```

```

resolution. "
374             "Re-run with -v for the native command + frame counts.",
375             video_id,
376             segmenter_actual,
377             detector_impl,
378         )
379     finally:
380         # Best-effort per-video observability report (never raises) ? parity with the
main CLI.
381         emit_indexer_report(
382             db=db,
383             video_id=video_id,
384             video_path=local_video,
385             config=config,
386             val_results=val_results,
387             val_context=val_context,
388             segmenter_requested=args.segmenter,
389             segmenter_actual=segmenter_actual,
390             was_reingest=False,
391             segmentation_ran=segmentation_ran,
392         )
393
394     # 5. SERIALIZE + 6. PUSH (outputs/<video_id>/?). idempotent on video_id.
395     pushed = _serialize_and_push_outputs(
396         scratch,
397         args.out_uri,
398         video_id,
399         segments,
400         status,
401         storage_cfg,
402         str(getattr(args, "video_uri", "") or ""),
403     )
404
405     print(
406         f"[worker] infer done: {len(segments)} segment(s); pushed {len(pushed)}
artifact(s):"
407     )
408     for u in pushed:
409         print("    ", u)
410
411     # M6: a remote dispatcher passes --done-uri; write the SUCCESS completion marker so
its bucket-poll
412     # terminates (carries video_id ? the dispatcher's output path + rc=0). Worker
FAILURES (rc 2/3)
413     # write their own marker via _fail() above, so the dispatcher fails FAST either way
(no 30-min hang
414     # on a bad video). DEFAULT-OFF: no --done-uri ? no marker written (the unchanged
local run).
415     if getattr(args, "done_uri", None):
416         _write_done_marker(
417             args.done_uri, video_id, 0, len(segments), status, storage_cfg, scratch
418         )
419     return 0
420
421
422     def _read_compile_spec(spec_uri: str, storage_cfg: dict, scratch: str) -> dict:
423         """Fetch + parse the control-plane-staged compile-spec (#521): the resolved
``time_pairs``, the
424         resolved ``stitching`` config, the source ``video_uri``, ``video_id``,
``output_name`` + ``locale``.
425         A gs://|s3:// URI is downloaded to scratch first (cloud backends require a real
dst)."""
426         local = storage.fetch(
427             spec_uri, os.path.join(scratch, "compile_spec.json"), config=storage_cfg
428         )

```

```

429         with open(local, encoding="utf-8") as f:
430             data = json.load(f)
431         if not isinstance(data, dict):
432             raise ValueError(f"compile spec at {spec_uri} is not a JSON object")
433         return data
434
435
436     def _write_compile_marker(
437         done_uri: str,
438         result: dict,
439         video_id: str,
440         returncode: int,
441         storage_cfg: dict,
442         scratch: str,
443     ) -> None:
444         """#521: write the compile done-marker (rc + status + the reel's
download_url/reel_uri) so the
445         dispatcher's bucket-poll terminates and can return the SPA result. Best-effort ?
never raises."""
446         try:
447             marker = {
448                 "video_id": video_id,
449                 "returncode": returncode,
450                 "status": result.get("status", "failed"),
451                 "download_url": result.get("download_url", ""),
452                 "reel_uri": result.get("reel_uri", ""),
453                 "message": result.get("message", ""),
454             }
455             local = os.path.join(scratch, "_compile_done.json")
456             with open(local, "w", encoding="utf-8") as f:
457                 json.dump(marker, f)
458             storage.put(local, done_uri, config=storage_cfg)
459             logger.info("[worker] wrote compile marker -> %s", done_uri)
460         except Exception as e: # never fail a good stitch because the marker push hiccuped
461             logger.warning("[worker] could not write compile done-marker %s: %s", done_uri,
e)
462
463
464     def cmd_compile(args: argparse.Namespace) -> int:
465         """#521: run the reel COMPILE/stitch on THIS worker (the L4 has NVENC + 8 vCPU). The
stitch is the
466         2026-07-08 CUJ's dominant term (~13 min on the 2-vCPU control-plane for an 11-clip
4K reel). Read
467         the control-plane-staged compile-spec (resolved time-pairs + stitching config +
source URI), stitch
468         the reel via the SAME ``orchestration._stitch_reel`` core the control-plane uses (so
there is no
469         forked stitch pipeline), upload the reel to the bucket, and write a done-marker for
the dispatcher's
470         bucket-poll. No control-plane DB is needed ? the spec carries everything the stitch
requires."""
471         from backend import orchestration
472         from backend.config import StitchingConfig
473
474         config = load_config(args.config) if args.config else load_config()
475         _apply_overrides(config, args.set)
476         storage_cfg = config.storage.model_dump()
477
478         scratch = args.scratch or tempfile.mkdtemp(prefix="rally-compile-")
479         os.makedirs(scratch, exist_ok=True)
480         # Clips + the assembled reel land in scratch; the reel is then mirror-uploaded to
the bucket.
481         config.output.output_dir = scratch
482         # The dispatcher forwards storage.output_root via --set; --out-uri carries the same
value as a

```



```

483         # fallback so reel_object_uri (which reads storage.output_root) always resolves the
reel's bucket
484         # destination even if the override were dropped.
485         if not (config.storage.output_root or "").strip():
486             config.storage.output_root = args.out_uri
487
488         spec = _read_compile_spec(args.spec_uri, storage_cfg, scratch)
489         video_id = str(spec.get("video_id") or "")
490         out_name = str(spec.get("output_name") or "")
491         source_ref = str(spec.get("video_uri") or "")
492         locale = spec.get("locale")
493         time_pairs = [
494             (float(pair[0]), float(pair[1])) for pair in (spec.get("time_pairs") or [])
495         ]
496         stitching = StitchingConfig(**(spec.get("stitching") or {}))
497
498         print(
499             f"[worker] compile: video_id={video_id} clips={len(time_pairs)} "
500             f"encoder={stitching.encoder} out={out_name}"
501         )
502
503         result = orchestration._stitch_reel(
504             config,
505             stitching,
506             video_id,
507             source_ref,
508             time_pairs,
509             out_name,
510             locale,
511             lambda stage: print(f"[worker] compile: {stage}"),
512         )
513         status = str(result.get("status") or "failed")
514         rc = 0 if status == "success" else 3
515         if status != "success":
516             print(f"[worker] compile FAILED: {result.get('message')}")
517
518         # DEFAULT-OFF: no --done-uri ? no marker (a direct-CLI stitch). A remote dispatcher
passes it.
519         if getattr(args, "done_uri", None):
520             _write_compile_marker(args.done_uri, result, video_id, rc, storage_cfg, scratch)
521         return rc
522
523
524     #: Video extensions accepted under <corpus>/videos/ (mirrors the labels/ '.csv' filter
so a
525     #: per-video sidecar ? .srt/.json/thumbnail ? can't shadow the real video on a stem
collision).
526     _VIDEO_EXTS = (".mp4", ".mov", ".mkv", ".avi", ".webm", ".m4v")
527
528     #: Leading ISO-date prefix on a canonical label name (`labels/<YYYY-MM-
DD>_<name>.rallies.csv`,
529     #: docs/DATA_IN_GCS.md §7.1). It disambiguates recycled GoPro names by authoring date
but is NOT
530     #: part of the clip's join key, so it's stripped when deriving the videos/ stem.
531     _LABEL_DATE_PREFIX = re.compile(r"^\d{4}-\d{2}-\d{2}_")
532
533
534     def _label_clip_name(fname: str) -> str:
535         """Derive the clip name (the ``videos/<name>.<ext>`` join key) from a label
filename.
536
537         Strips both the ``.rallies.csv`` / ``.csv`` suffix and an optional leading ISO-date
prefix, so a
538         canonical dated label and a legacy bare label resolve to the SAME bare clip stem:
539

```

```

540      * ``2026-06-22_GX020094.rallies.csv`` -> ``GX020094`` (canonical, date-
stamped)
541      * ``GX020094.rallies.csv`` -> ``GX020094`` (legacy bare ? same
stem)
542      * ``GX020094_proxy.rallies.csv`` -> ``GX020094_proxy`` (``_proxy`` is part of
the name, kept)
543      * ``2026-06-21_mahadevpura_1.rallies.csv`` -> ``mahadevpura_1`` (underscores in
<name> survive)
544
545      Without the date strip, ``--trajectory-source detector`` can't match a dated label
to its bare
546      video stem, so every clip is skipped and train aborts with ``<2 videos``
(DATA_IN_GCS $7b #1).
547      """
548      name = fname.replace(".rallies.csv", "").replace(".csv", "")
549      return _LABEL_DATE_PREFIX.sub("", name)
550
551
552      def _probe_video_fps_width(path: str):
553          """Robust ffprobe (fps, frame_width) for a real video. Returns None on ANY failure.
554
555          The detector trajectory is frame-indexed; the harness maps frame->time via fps, so a
WRONG
556          fps silently mis-aligns every rally label (e.g. a 60fps clip read as 30 doubles
every time).
557          cv2's ``_probe_fps_width`` silently defaults to (30, 1280) on a read miss ? fine for
the
558          inference report (it flags the fallback) but corrupting for REAL training data. So
here we
559          probe with ffprobe (same tool as ``get_video_duration``) and return None so the
caller SKIPS
560          rather than guesses.
561          """
562          import json
563          import subprocess
564
565          try:
566              out = subprocess.check_output(
567                  [
568                      ffprobe_bin(),
569                      "-v",
570                      "error",
571                      "-select_streams",
572                      "v:0",
573                      "-show_entries",
574                      "stream=r_frame_rate,width",
575                      "-of",
576                      "json",
577                      path,
578                  ],
579                  stderr=subprocess.DEVNULL,
580              ).decode()
581              st = (json.loads(out).get("streams") or [{}])[0]
582              num, _, den = str(st.get("r_frame_rate", "0/1")).partition("/")
583              fps = float(num) / float(den or "1")
584              width = float(st.get("width") or 0)
585              if fps > 0 and width > 0:
586                  return fps, width
587          except (
588              subprocess.SubprocessError,
589              ValueError,
590              KeyError,
591              OSError,
592              json.JSONDecodeError,
593          ):

```

```

594         pass
595     return None
596
597
598 def _resolve_train_detector(args: argparse.Namespace, config: Any):
599     """Build the trajectory DetectorRunner for `--trajectory-source detector`, honouring
600     detector_impl (RALLY_DETECTOR_IMPL / indexing.detector_impl: auto=WSL, native=GPU
box,
601     stub=CPU mock). Returns the runner, or prints an error + returns None.
602
603     Reuses the segmenter's `_resolve_runner` seam so the worker and the live CLI build
the
604     detector identically (one source of truth). A non-trajectory segmenter (e.g. yolo)
has
605     no shuttle detector and is rejected.
606     """
607     from backend.pipeline.segmenters.trajectory_hybrid import TrajectoryHybridSegmenter
608
609     seg_cls = segmenter_registry.get(args.segmenter)
610     if not seg_cls:
611         print(
612             f"[worker] unknown segmenter: {args.segmenter!r} "
613             f"(have {list(segmenter_registry.list_available())})"
614         )
615         return None
616     seg = seg_cls(None, config)
617     if not isinstance(seg, TrajectoryHybridSegmenter):
618         print(
619             f"[worker] --segmenter {args.segmenter!r} has no shuttle detector; "
620             f"use a trajectory hybrid (wasb_hybrid / tracknet_hybrid / fusion_hybrid)."
621         )
622         return None
623     try:
624         runner, _ = seg._resolve_runner(config.indexing)
625     except NotImplementedError as e:
626         # e.g. detector_impl='native' for a family without a native runner (tracknet, or
627         # fusion with a tracknet substrate). Fail cleanly (rc!=0), not with a raw
traceback.
628         print(f"[worker] detector not available: {e}")
629         return None
630     ok, msg = runner.healthcheck()
631     if not ok:
632         print(f"[worker] detector environment not ready: {msg}")
633         return None
634     print(f"[worker] detector ready ({args.segmenter}): {msg}")
635     return runner
636
637
638 def cmd_train(args: argparse.Namespace) -> int:
639     """Gen-0 train-from-bucket: pull labels (+ videos) -> trajectories -> manifest ->
shell out
640     to training.gen0.harness (backend must NOT import training) -> push the artifact
bundle.
641
642     Two trajectory sources (the train pipeline + harness are identical either way):
643     * ``synth`` (default) ? the CPU "mock GPU": deterministic, label-consistent
synthetic
644         shuttle paths (no GPU/WSL), so the whole pipeline is validated on a CPU box /
CI.
645     * ``detector`` ? REAL shuttle trajectories: pull each video and run the resolved
646         DetectorRunner (native WASB on a GPU box, or stub for a GPU-free wiring test).
This
647         is the M3.5 "first real training" path; only the trajectory source flips.
648     """
649     import subprocess

```

```

650
651 config = load_config(args.config) if args.config else load_config()
652 _apply_overrides(config, args.set)
653 if not _run_startup_checks(config):
654     return 2
655 storage_cfg = config.storage.model_dump()
656
657 scratch = args.scratch or tempfile.mkdtemp(prefix="rally-train-")
658 labels_dir = os.path.join(scratch, "labels")
659 traj_dir = os.path.join(scratch, "trajectories")
660 videos_dir = os.path.join(scratch, "videos")
661 os.makedirs(labels_dir, exist_ok=True)
662 os.makedirs(traj_dir, exist_ok=True)
663
664 corpus = args.corpus_uri.rstrip("/")
665 label_uris = sorted(
666     u
667     for u in storage.list_uris(f"{corpus}/labels/", config=storage_cfg)
668     if u.lower().endswith(".csv")
669 )
670 if len(label_uris) < 2:
671     print(
672         f"[worker] need >=2 labeled videos for leave-one-video-out; found
{len(label_uris)} "
673         f"under {corpus}/labels/"
674     )
675     return 2
676
677 # Real-detector source: resolve the runner once + index the bucket's videos/ by
stem.
678 runner = None
679 video_by_stem: dict = {}
680 if args.trajectory_source == "detector":
681     os.makedirs(videos_dir, exist_ok=True)
682     runner = _resolve_train_detector(args, config)
683     if runner is None:
684         return 2
685     for vuri in storage.list_uris(f"{corpus}/videos/", config=storage_cfg):
686         vname = storage.parse_uri(vuri).name
687         if not vname.lower().endswith(_VIDEO_EXTS):
688             continue # skip sidecars (.srt/.json/thumbnails) so they can't shadow
the video
689         video_by_stem[os.path.splitext(vname)[0]] = vuri
690
691 print(f"[worker] trajectory source: {args.trajectory_source}")
692 specs: List[dict] = []
693 for uri in label_uris:
694     fname = storage.parse_uri(uri).name # e.g. 2026-06-22_GX020094.rallies.csv
695     name = _label_clip_name(fname) # -> GX020094 (date prefix stripped)
696     local_label = storage.fetch(
697         uri, os.path.join(labels_dir, fname), config=storage_cfg
698     )
699     if args.trajectory_source == "detector":
700         vuri = video_by_stem.get(name)
701         if not vuri:
702             print(
703                 f"[worker] no video for {name!r} under {corpus}/videos/ ? skipping."
704             )
705             continue
706         local_video = storage.fetch(
707             vuri,
708             os.path.join(videos_dir, storage.parse_uri(vuri).name),
709             config=storage_cfg,
710         )
711         video_id = compute_video_id(local_video)

```

```

712         # Real fps/width drive the trajectory frame->time mapping; PROBE (don't
guess) ? a
713         # wrong fps silently mis-aligns every label, so skip the video if ffprobe
can't read it.
714         meta = _probe_video_fps_width(local_video)
715         if meta is None:
716             print(
717                 f"[worker] could not probe fps/width for {name!r} (ffprobe) ?
skipping (won't guess).")
718             )
719             continue
720         fps, frame_width = meta
721         traj = runner.run_predict(local_video, traj_dir, video_id=video_id)
722         if not traj:
723             print(
724                 f"[worker] detector produced no trajectory for {name!r} ? skipping."
725             )
726             continue
727         specs.append(
728             {
729                 "name": name,
730                 "trajectory": traj,
731                 "labels": local_label,
732                 "fps": fps,
733                 "frame_width": frame_width,
734             }
735         )
736     else:
737         from backend.pipeline.detectors.stub_runner import (
738             synth_trajectory_from_labels,
739         )
740
741         traj = os.path.join(traj_dir, f"{name}_traj.csv")
742         synth_trajectory_from_labels(
743             local_label, traj, fps=args.fps, frame_width=args.frame_width
744         )
745         specs.append(
746             {
747                 "name": name,
748                 "trajectory": traj,
749                 "labels": local_label,
750                 "fps": args.fps,
751                 "frame_width": args.frame_width,
752             }
753         )
754
755     if len(specs) < 2:
756         print(
757             f"[worker] need >=2 videos with trajectories for leave-one-video-out; got
{len(specs)}."
758         )
759         return 2
760     manifest_path = os.path.join(scratch, "manifest.json")
761     with open(manifest_path, "w", encoding="utf-8") as f:
762         json.dump(specs, f, indent=2)
763     src_label = (
764         "real detector" if args.trajectory_source == "detector" else "CPU-mock stub"
765     )
766     print(f"[worker] built manifest: {len(specs)} videos ({src_label} trajectories)")
767
768     out_dir = os.path.join(scratch, "artifacts")
769     cmd = [
770         sys.executable,
771         "-m",
772         "training.gen0.harness",

```

```

773         "--manifest",
774         manifest_path,
775         "--out",
776         out_dir,
777         "--version",
778         args.version,
779     ]
780     if args.floor:
781         floor_local = (
782             storage.fetch(
783                 args.floor, os.path.join(scratch, "floor.json"), config=storage_cfg
784             )
785             if "://" in args.floor
786             else args.floor
787         )
788         cmd += ["--floor", floor_local]
789     print("[worker] training:", " ".join(cmd))
790     rc = subprocess.run(cmd).returncode
791     if rc != 0:
792         print(f"[worker] gen0 harness failed (rc={rc})")
793         return rc
794
795     # PUSH the versioned artifact bundle (weights.json + manifest.json) to the bucket.
796     bundle = os.path.join(out_dir, args.version)
797     out_prefix = args.out_uri.rstrip("/")
798     pushed = []
799     for fn in sorted(os.listdir(bundle)):
800         uri = f"{out_prefix}/weights/{args.version}/{fn}"
801         storage.put(os.path.join(bundle, fn), uri, config=storage_cfg)
802         pushed.append(uri)
803     print(f"[worker] train done: pushed {len(pushed)} artifact(s):")
804     for u in pushed:
805         print("    ", u)
806     return 0
807
808
809 def build_parser() -> argparse.ArgumentParser:
810     p = argparse.ArgumentParser(
811         prog="python -m backend.worker",
812         description="Storage-aware worker: pull ? run pipeline ? push.",
813     )
814     p.add_argument(
815         "-v",
816         "--verbose",
817         action="store_true",
818         help="DEBUG logging (the native detector command, frame counts, per-stage detail)",
819     )
820     sub = p.add_subparsers(dest="cmd", required=True)
821
822     pi = sub.add_parser("infer", help="run inference on one video URI and push outputs")
823     pi.add_argument(
824         "--video-uri", required=True, help="local path / local:// / s3:// / gcs://"
825     )
826     pi.add_argument(
827         "--out-uri",
828         required=True,
829         help="output prefix (outputs/<video_id>/? is appended)",
830     )
831     pi.add_argument(
832         "--segmenter", default=None, help="default: indexing.default_segmenter"
833     )
834     pi.add_argument(
835         "--config", default=None, help="config.json path (default: ./config.json)"
836     )

```

```

837     pi.add_argument(
838         "--scratch", default=None, help="local scratch dir (default: a temp dir)"
839     )
840     pi.add_argument("--max-frames", type=int, default=None)
841     pi.add_argument(
842         "--done-uri",
843         default=None,
844         help="M6: storage URI for a completion MARKER (video_id + rc) written on a "
845         "successful run, so a remote (Cloud Run) dispatcher's bucket-poll terminates. "
846         "Default-OFF: omit it for the normal local run.",
847     )
848     pi.add_argument(
849         "--set",
850         action="append",
851         default=[],
852         metavar="k=v",
853         help="config override, e.g. --set indexing.skip_ai_handoff=true",
854     )
855     pi.add_argument(
856         "--skip-validation",
857         action="store_true",
858         help="skip the worker's re-validation: the dispatcher (control-plane) already
ran the "
859         "validators with its resolved config and only dispatches on PASS, so re-
validating here is "
860         "redundant, decodes the video again, and can contradict the gate. Default-OFF:
the direct "
861         "`worker infer` CLI still validates.",
862     )
863     pi.set_defaults(func=cmd_infer)
864
865     pc = sub.add_parser(
866         "compile",
867         help="#521: stitch a highlight reel on this worker (L4 NVENC) from a staged
compile-spec",
868     )
869     pc.add_argument(
870         "--spec-uri",
871         required=True,
872         help="storage URI of the control-plane-staged compile-spec JSON (time-pairs +
stitching "
873         "config + source uri + video_id)",
874     )
875     pc.add_argument(
876         "--out-uri",
877         default="",
878         help="output ROOT (gs://? bucket); the reel is mirrored to
highlights/<video_id>/<name>. "
879         "Also forwarded via --set storage.output_root; either resolves the reel path.",
880     )
881     pc.add_argument(
882         "--done-uri",
883         default=None,
884         help="storage URI for the compile completion MARKER (rc + download_url), so a
remote "
885         "dispatcher's bucket-poll terminates. Default-OFF: omit for a direct-CLI
stitch.",
886     )
887     pc.add_argument(
888         "--config", default=None, help="config.json path (default: ./config.json)"
889     )
890     pc.add_argument(
891         "--scratch", default=None, help="local scratch dir (default: a temp dir)"
892     )
893     pc.add_argument(

```

```

894         "--set",
895         action="append",
896         default=[],
897         metavar="k=v",
898         help="config override, e.g. --set storage.output_root=gs://bucket",
899     )
900     pc.set_defaults(func=cmd_compile)
901
902     pt = sub.add_parser(
903         "train",
904         help="Gen-0 train-from-bucket (labels [+videos] -> trajectories -> harness)",
905     )
906     pt.add_argument(
907         "--corpus-uri",
908         required=True,
909         help="prefix containing labels/ (+ videos/ for "
910         "--trajectory-source detector), e.g. s3://bucket or gcs://bucket",
911     )
912     pt.add_argument(
913         "--out-uri",
914         required=True,
915         help="output prefix (weights/<version>/ is appended)",
916     )
917     pt.add_argument("--version", required=True, help="semver for the artifact bundle")
918     pt.add_argument(
919         "--trajectory-source",
920         choices=["synth", "detector"],
921         default="synth",
922         help="synth = CPU-mock label-consistent trajectories (default, no GPU); "
923         "detector = REAL shuttle trajectories via the resolved DetectorRunner "
924         "(needs videos/ in the corpus; honours detector_impl).",
925     )
926     pt.add_argument(
927         "--segmenter",
928         default="wasb_hybrid",
929         help="trajectory hybrid whose detector is used for --trajectory-source
930         detector",
931     )
932     pt.add_argument(
933         "--floor", default=None, help="floor baseline JSON (local path or storage URI)"
934     )
935     pt.add_argument("--config", default=None)
936     pt.add_argument(
937         "--scratch", default=None
938     )
939     pt.add_argument(
940         "--fps",
941         type=float,
942         default=30.0,
943         help="synth trajectory fps (detector probes the video)",
944     )
945     pt.add_argument(
946         "--frame-width",
947         type=float,
948         default=1280.0,
949         help="synth trajectory frame width (detector probes the video)",
950     )
951     pt.add_argument("--set", action="append", default=[], metavar="k=v")
952     pt.set_defaults(func=cmd_train)
953     return p
954
955     def _setup_logging(verbose: bool) -> None:
956         """Configure logging for the worker CLI so the pipeline's existing INFO/ERROR
957         diagnostics are
958         VISIBLE ? the worker previously ran with no logging config, so every
959         segmenter/detector log

```



```

956         (detector choice, 'Running native WASB inference', availability failures) was
silently dropped.
957         That is exactly why a 0-rally cold-start gave nothing to analyse. Configured here
(the CLI
958         entry), not at import, so importing the module never reconfigures a caller's
logging.""
959         # force=True: basicConfig is a no-op if the root already has handlers (e.g. a
library
960         # configured logging at import), which would leave the level at WARNING and re-
suppress the
961         # diagnostics. Forcing reconfiguration means the worker reliably controls its own
verbosity.
962         logging.basicConfig(
963             level=logging.DEBUG if verbose else logging.INFO,
964             format="%(asctime)s [%(levelname)s] %(name)s: %(message)s",
965             datefmt="%H:%M:%S",
966             force=True,
967         )
968
969
970     def main(argv=None) -> int:
971         args = build_parser().parse_args(argv)
972         _setup_logging(getattr(args, "verbose", False))
973         return args.func(args)
974
975
976     if __name__ == "__main__":
977         sys.exit(main())
978

```

13. assistant (2026-07-08T06:15:53.478Z)

Now I've confirmed the worker side: `cmd\_compile` (worker/`\_\_main\_\_.py`:513-514) sets `rc = 0` if status == "success" purely from `\_stitch\_reel`'s return, with no post-mirror verification. Let me read the rest of `\_dispatch\_compile\_remote` (control-plane side ? does IT verify the reel exists?), and check the test coverage.

14. user (2026-07-08T06:15:56.281Z)

```

"""#521: run the compile/stitch on the L4 worker instead of the CPU-only control-plane. The
control-plane owns the DB, so it RESOLVES the reel's segments here (`_resolve_compile`),
stages a
small compile-spec (resolved time-pairs + the resolved stitching/watermark config + the
source URI)
to the bucket, dispatches a ``compile`` Cloud Run Job on the L4 (NVENC), bucket-polls the
done-marker, and returns the same ``{status, filepath, download_url}`` shape the SPA
expects. The
worker needs NO control-plane DB ? everything it needs rides the spec. Mirrors
``_maybe_dispatch_remote`` (the DETECT dispatch)."""
from backend.compute import CompileJobSpec, get_compute_backend

notify("resolving")
try:
    video, kept, out_name = _resolve_compile(
        db, cfg, video_id, segment_ids, output_filename
    )
except _CompileError as e:
    return {"status": "failed", "message": e.detail, "video_id": video_id}

stitching = getattr(cfg, "stitching", None) or StitchingConfig()
storage_cfg = getattr(cfg, "storage", None)
out_root = (getattr(storage_cfg, "output_root", "") or "").strip()
compute = get_compute_backend(cfg, target_override="cloud_run") if out_root else None
if compute is None:
    # Compile was placed on the L4 but the cloud dispatch isn't actually wired here (no

```

```

output
    # bucket / Cloud Run env). Fall back to the in-process stitch rather than fail ? a slow
reel
    # beats no reel. (The placement guard means this is only reachable in a mis-provisioned
setup.)
    logger.warning(
        "[compile] compile placed on the L4 but no cloud backend available (out_root=%r) ? "
        "stitching in-process on the control-plane instead.",
        out_root,
    )
    time_pairs = [(s["start_time"], s["end_time"]) for s in kept]
    return _stitch_reel(
        cfg, stitching, video_id, video["filepath"], time_pairs, out_name, locale, notify
    )

    # Build the compile-spec the worker consumes. time-pairs are JSON lists (frozen-safe on the
wire);
    # the stitching config is the control-plane's RESOLVED one, with the encoder swapped to the
L4's
    # (deployment.cloud_stitch_encoder, NVENC by default) so the worker hardware-encodes the
reel.
    stitch_dump = stitching.model_dump()
    stitch_dump["encoder"] = getattr(
        getattr(cfg, "deployment", None), "cloud_stitch_encoder", "h264_nvenc"
    )
    spec_payload = {
        "video_id": video_id,
        "video_uri": video["filepath"],
        "output_name": out_name,
        "locale": locale,
        "time_pairs": [
            [float(s["start_time"]), float(s["end_time"])] for s in kept
        ],
        "stitching": stitch_dump,
    }
    token = uuid.uuid4().hex
    spec_uri = f"{out_root.rstrip('/')}/_compile/{token}/spec.json"
    try:
        notify("staging compile spec")
        with tempfile.TemporaryDirectory(prefix="cr-compile-") as td:
            local_spec = os.path.join(td, "spec.json")
            with open(local_spec, "w", encoding="utf-8") as f:
                json.dump(spec_payload, f)
            storage.put(local_spec, spec_uri, config=_storage_settings(cfg))
    except Exception as e: # noqa: BLE001 ? a staging failure is a clean job failure, not a
crash
        logger.error(
            "[compile] failed to stage the compile spec to %s: %s", spec_uri, e, exc_info=True
        )
        return {
            "status": "failed",
            "message": f"cloud compile: could not stage the compile spec: {e}",
            "video_id": video_id,
        }

    # Forward the bucket coordinates so the worker (no DEPLOYMENT_PROFILE) can read the spec,
mirror the
    # reel, and write the marker ? mirrors how _maybe_dispatch_remote forwards the WASB knobs
for detect.
    overrides = [f"storage.output_root={out_root}"]
    backend_name = (getattr(storage_cfg, "backend", "") or "").strip()
    if backend_name:
        overrides.append(f"storage.backend={backend_name}")
    spec = CompileJobSpec(
        out_uri=out_root,

```

```

        spec_uri=spec_uri,
        video_id=video_id,
        output_name=out_name,
        config_overrides=tuple(overrides),
    )

    notify("dispatching stitch (cloud_run)")
    logger.info(
        "[compile] dispatching stitch for video_id=%s -> Cloud Run Job (out=%s, %d clip(s))",
        video_id,
        out_root,
        len(kept),
    )
    try:
        result = compute.run_compile(
            spec, on_wait=lambda elapsed: notify(_compile_wait_message(elapsed))
        )
    except Exception as e: # noqa: BLE001 ? any dispatch/poll error is a clean job failure, not
a 500
        logger.error("[compile] cloud compile dispatch error: %s", e, exc_info=True)
        return {
            "status": "failed",
            "message": f"cloud compile dispatch error: {e}",
            "video_id": video_id,
        }
    if result.returncode != 0:
        msg = str(
            result.report.get("message")
            or f"cloud compile job failed (returncode {result.returncode})."
        )
        return {"status": "failed", "message": msg, "video_id": video_id}

    download_url = str(
        result.report.get("download_url")
        or f"/api/highlights/{quote(video_id)}/{quote(out_name)}"
    )
    logger.info(
        "[COMPILE] cloud_run stitched: execution=%s out=%s video_id=%s",
        result.execution,
        result.outputs_uri,
        video_id,
    )
    notify("finished")
    return {
        "status": "success",
        "filepath": str(result.report.get("reel_uri") or result.outputs_uri or ""),
        "download_url": download_url,
        "video_id": video_id,
        "compute": {
            "path": "cloud_run",
            "execution": result.execution,
            "job": os.environ.get("CLOUD_RUN_JOB", "rally-worker"),
            "region": os.environ.get("CLOUD_RUN_REGION", "asia-south1"),
        },
    }
}

```

15. user (2026-07-08T06:15:57.206Z)

```

952: def test_parser_compile_accepts_spec_and_done_uris():
989: def test_cmd_compile_stitches_and_writes_marker(tmp_path, monkeypatch):
990:     """cmd_compile reads the staged spec, runs the SHARED _stitch_reel core (VideoCompiler
stubbed so
991:     no ffmpeg), mirrors the reel, and writes a compile done-marker for the dispatcher's
bucket-poll."""
1012:         "reel_uri": f"gs://b/highlights/{video_id}/{out_name}",

```

```

1015: monkeypatch.setattr(orch, "_stitch_reel", _fake_stitch)
1048: def test_cmd_compile_failure_returns_rc3_and_failed_marker(tmp_path, monkeypatch):
1062:     "_stitch_reel",
1089: def test_cmd_compile_no_done_uri_writes_no_marker(tmp_path, monkeypatch):
1100:     orch, "_stitch_reel", lambda *a, **k: {"status": "success", "video_id": "vidReel"}

```

16. assistant (2026-07-08T06:16:11.540Z)

Now let me verify the full test-coverage picture: search ALL test files for any test that exercises the real `\_stitch\_reel` (not mocked) and specifically a mirror/`storage.put` failure, plus check `test\_compile\_dispatch.py` / `test\_compute\_backend.py` for the fake client.

17. user (2026-07-08T06:16:13.407Z)

```

tests\test_compile_dispatch.py:174:     """DEFAULT (compile on the control-plane):
_execute_compile stitches in-process via _stitch_reel
tests\test_compile_dispatch.py:188:     monkeypatch.setattr(orch, "_stitch_reel", _fake_stitch)
tests\test_worker.py:990:     """cmd_compile reads the staged spec, runs the SHARED _stitch_reel
core (VideoCompiler stubbed so
tests\test_worker.py:1015:     monkeypatch.setattr(orch, "_stitch_reel", _fake_stitch)
tests\test_worker.py:1062:     "_stitch_reel",
tests\test_worker.py:1100:     orch, "_stitch_reel", lambda *a, **k: {"status": "success",
"video_id": "vidReel"}

```

18. user (2026-07-08T06:16:13.787Z)

```

985         json.dump(spec, f)
986     return spec_uri
987
988
989     def test_cmd_compile_stitches_and_writes_marker(tmp_path, monkeypatch):
990         """cmd_compile reads the staged spec, runs the SHARED _stitch_reel core
(VideoCompiler stubbed so
991         no ffmpeg), mirrors the reel, and writes a compile done-marker for the dispatcher's
bucket-poll."""
992         import backend.orchestration as orch
993
994         bucket = tmp_path / "bucket"
995         bucket.mkdir()
996         source = tmp_path / "src.mp4"
997         source.write_bytes(b"SRC")
998         spec_uri = _write_compile_spec(bucket, "vidReel", source, "reel.mp4")
999         done = str(bucket / "_compile" / "tok.done")
1000
1001         seen = {}
1002
1003         def _fake_stitch(cfg, stitching, video_id, source_ref, time_pairs, out_name, locale,
notify):
1004             seen["encoder"] = stitching.encoder
1005             seen["pairs"] = time_pairs
1006             seen["source"] = source_ref
1007             notify("stitching")
1008             return {
1009                 "status": "success",
1010                 "video_id": video_id,
1011                 "download_url": f"/api/highlights/{video_id}/{out_name}",
1012                 "reel_uri": f"gs://b/highlights/{video_id}/{out_name}",
1013             }
1014
1015         monkeypatch.setattr(orch, "_stitch_reel", _fake_stitch)
1016
1017         cfg = tmp_path / "c.json"
1018         cfg.write_text('{"storage": {"backend": "local"}}')
1019         rc = wmain.cmd_compile(

```

```

1020         wmain.build_parser().parse_args(
1021             [
1022                 "compile",
1023                 "--spec-uri",
1024                 spec_uri,
1025                 "--out-uri",
1026                 str(bucket),
1027                 "--done-uri",
1028                 done,
1029                 "--config",
1030                 str(cfg),
1031                 "--scratch",
1032                 str(tmp_path / "s"),
1033                 "--set",
1034                 f"storage.output_root={bucket}",
1035             ]
1036         )
1037     )
1038     assert rc == 0
1039     # the shared stitch core got the resolved pairs + the NVENC encoder from the spec
1040     assert seen["encoder"] == "h264_nvenc"
1041     assert seen["pairs"] == [(2.0, 6.0), (10.0, 14.0)]
1042     assert seen["source"] == str(source)
1043     marker = json.loads(open(done, encoding="utf-8").read())
1044     assert marker["returncode"] == 0 and marker["status"] == "success"
1045     assert marker["download_url"] == "/api/highlights/vidReel/reel.mp4"
1046
1047
1048 def test_cmd_compile_failure_returns_rc3_and_failed_marker(tmp_path, monkeypatch):
1049     """A stitch failure surfaces as rc 3 + a failed marker (rc != 0) so the dispatcher's
1050     poll
1051     terminates and the SPA sees a clean failure."""
1052     import backend.orchestration as orch
1053
1054     bucket = tmp_path / "bucket"
1055     bucket.mkdir()
1056     source = tmp_path / "src.mp4"
1057     source.write_bytes(b"SRC")
1058     spec_uri = _write_compile_spec(bucket, "vidReel", source, "reel.mp4")
1059     done = str(bucket / "_compile" / "tok.done")
1060
1061     monkeypatch.setattr(
1062         orch,
1063         "_stitch_reel",
1064         lambda *a, **k: {"status": "failed", "message": "ffmpeg blew up", "video_id":
1065         "vidReel"},
1066     )
1067     cfg = tmp_path / "c.json"
1068     cfg.write_text('{"storage": {"backend": "local"}}')
1069     rc = wmain.cmd_compile(
1070         wmain.build_parser().parse_args(
1071             [
1072                 "compile",
1073                 "--spec-uri",
1074                 spec_uri,
1075                 "--out-uri",
1076                 str(bucket),
1077                 "--done-uri",
1078                 done,
1079                 "--config",
1080                 str(cfg),
1081                 "--scratch",
1082                 str(tmp_path / "s"),

```

```

1083     )
1084     assert rc == 3
1085     marker = json.loads(open(done, encoding="utf-8").read())
1086     assert marker["returncode"] == 3 and marker["status"] == "failed"
1087
1088
1089     def test_cmd_compile_no_done_uri_writes_no_marker(tmp_path, monkeypatch):
1090         """DEFAULT-OFF: a direct-CLI stitch with no --done-uri writes no marker."""
1091         import backend.orchestration as orch
1092
1093         bucket = tmp_path / "bucket"
1094         bucket.mkdir()
1095         source = tmp_path / "src.mp4"
1096         source.write_bytes(b"SRC")
1097         spec_uri = _write_compile_spec(bucket, "vidReel", source, "reel.mp4")
1098
1099         monkeypatch.setattr(

```

19. assistant (2026-07-08T06:17:01.575Z)

I've confirmed all 4 `\_stitch\_reel` references in tests are monkeypatches. But a test could invoke the real `\_stitch\_reel` transitively via `cmd\_compile` without mentioning it by name. Let me verify `test\_compute\_backend.py` and `test\_compile\_dispatch.py` don't drive the real worker mirror path with an injected `storage.put` failure.

20. user (2026-07-08T06:17:05.897Z)

```

160         by the control-plane first."""
161         db, cfg, video_id, seg_ids, _bucket, monkeypatch = compile_env
162         fake = _FakeCompileClient()
163         _inject_fake(monkeypatch, fake)
164
165         out = orch._execute_compile(
166             cfg, db, _job(video_id, seg_ids, "myreel-sd.mp4", rendition="sd")
167         )
168
169         assert out["status"] == "success"
170         assert out["filename"] == "myreel-sd.mp4"
171         assert out["rendition"] == "sd"
172         assert fake.spec is not None
173         assert fake.spec["output_name"] == "myreel-sd.mp4"
174         assert fake.spec["rendition"] == "sd"
175         assert fake.spec["stitching"]["renditions"]["sd"]["max_width"] == 854
176         assert fake.spec["stitching"]["renditions"]["sd"]["max_height"] == 480
177         assert len(fake.calls) == 1
178
179
180     def test_compile_dispatch_forwards_stitch_config_overrides(compile_env):
181         """The worker Job has no DEPLOYMENT_PROFILE, so the dispatch forwards the bucket
coordinates
182         (storage.output_root/backend) ? else the worker can't read the spec or mirror the
reel."""
183         db, cfg, video_id, seg_ids, bucket, monkeypatch = compile_env
184         fake = _FakeCompileClient()
185         _inject_fake(monkeypatch, fake)
186         orch._execute_compile(cfg, db, _job(video_id, seg_ids))
187         joined = " ".join(fake.calls[0])
188         assert f"storage.output_root={bucket}" in joined
189         assert "storage.backend=local" in joined
190
191
192     def test_compile_cloud_failure_returns_failed(compile_env):
193         """A worker stitch failure (rc != 0) surfaces as a failed job with the worker's
message ? the SPA
194         gets a clean verdict, not a crash."""

```

```

195     db, cfg, video_id, seg_ids, _bucket, monkeypatch = compile_env
196     fake = _FakeCompileClient(returncode=3)
197     _inject_fake(monkeypatch, fake)
198
199     out = orch._execute_compile(cfg, db, _job(video_id, seg_ids))
200     assert out["status"] == "failed"
201     assert "stitch blew up" in out["message"]
202     assert out["video_id"] == video_id
203
204
205     def test_compile_default_stitches_in_process_no_dispatch(compile_env, monkeypatch):
206         """DEFAULT (compile on the control-plane): _execute_compile stitches in-process via
207         _stitch_reel
208         and NEVER dispatches to Cloud Run. The stitch core is stubbed so no real ffmpeg
209         runs."""
210         db, cfg, video_id, seg_ids, _bucket, _mp = compile_env
211         # flip compile back to the control-plane (the default) ? everything else stays
212         cloud-wired
213         cfg.deployment.phase_placement.compile = "control_plane"
214         fake = _FakeCompileClient()
215         _inject_fake(monkeypatch, fake)
216         seen = {}
217
218         def _fake_stitch(c, stitching, vid, source, pairs, name, locale, notify,
219             rendition="original"):
220             seen["pairs"] = pairs
221             seen["source"] = source
222             seen["rendition"] = rendition
223             return {"status": "success", "video_id": vid, "download_url":
224                 f"/api/highlights/{vid}/{name}"}
225
226         monkeypatch.setattr(orch, "_stitch_reel", _fake_stitch)
227
228         out = orch._execute_compile(cfg, db, _job(video_id, seg_ids, "local.mp4"))
229         assert out["status"] == "success"
230         assert fake.calls == [] # NEVER dispatched ? the control-plane stitched in-process
231         assert seen["pairs"] == [(2.0, 6.0), (10.0, 14.0), (20.0, 22.0)]
232         assert seen["source"].endswith("source.mp4")
233         assert seen["rendition"] == "original"

```

21. user (2026-07-08T06:17:06.378Z)

```

=== test_compute_backend: compile/CloudRunJobClient refs ===
3:All GPU-free + cloud-free: the Cloud Run trigger is a FAKE ``CloudRunJobClient`` that
simulates the
4:worker (writes the done-marker + report to a local "bucket" = a tmp dir), so the
dispatch?bucket-poll
20:from backend.compute.cloud_run import CloudRunCompute, CloudRunJobClient
24:# A fast, no-wait policy (the fake client writes the marker synchronously, so the first poll
hits).
28:class _FakeRunClient(CloudRunJobClient):
29:    """Simulates a Cloud Run Job: on dispatch, write the report + done-marker into the local
bucket
76:class _NoopRunClient(CloudRunJobClient):
94:class _NeverDoneClient(CloudRunJobClient):
117:class _CrashedNoMarkerClient(CloudRunJobClient):
176:    """The real GoogleCloudRunJobClient must fail loudly when GOOGLE_CLOUD_PROJECT is unset
?
180:    from backend.compute.cloud_run import GoogleCloudRunJobClient
182:    client = GoogleCloudRunJobClient()
195:    from backend.compute.cloud_run import GoogleCloudRunJobClient
197:    client = GoogleCloudRunJobClient()
598:    from backend.compute.cloud_run import GoogleCloudRunJobClient
600:    client = GoogleCloudRunJobClient()

```

```

764: def test_worker_writes_done_marker_on_success(tmp_path, monkeypatch):
765:     """The container side: cmd_infer with --done-uri writes a completion marker (video_id +
rc) so a
834:     """DEFAULT-OFF: without --done-uri the inline run writes NO marker (today's
behaviour)."""
888: def test_worker_writes_failure_marker_on_validation_fail(tmp_path, monkeypatch):
889:     """The container side: a worker-side FAILURE (validation rc=2) with --done-uri writes a
marker
948: def test_worker_failure_writes_no_marker_without_done_uri(tmp_path, monkeypatch):
949:     """DEFAULT-OFF: a worker failure with NO --done-uri writes no marker (the unchanged
local run)."""
1022:     # The cloud worker reads the gs:// input + writes the gs:// outputs/done-marker through
1084:     client = _FakeRunClient() # still writes report.json so the post-marker read succeeds
1108: class _FakeCompileClient(CloudRunJobClient):
1109:     """Simulates the Cloud Run COMPILE worker: on dispatch, write a reel done-marker into
the local
1110:     bucket exactly where the real ``cmd_compile`` would, so ``run_compile``'s bucket-poll
resolves."""
1137: def test_run_compile_dispatches_and_returns_marker(tmp_path):
1138:     """run_compile dispatches a `compile` subcommand (spec + done marker), bucket-polls,
and returns
1147:     result = backend.run_compile(
1168: def test_run_compile_surfaces_nonzero_returncode(tmp_path):
1178:     result = backend.run_compile(
1184: def test_compile_marker_prefix_distinct_from_infer(tmp_path):
1194:     backend.run_compile(
1201: def test_run_compile_default_backend_raises_not_implemented():
1203:     CloudRunCompute overrides run_compile) ? never a silent no-op."""
1211:     _InProc().run_compile(

```

22. assistant (2026-07-08T06:17:51.331Z)

I have now fully confirmed both halves of the claim. Let me consolidate the findings.

```

**Failure mode is real (confirmed in code):**
- `backend/orchestration.py` `_stitch_reel` (HEAD): after a successful local stitch,
`storage.put(out_path, reel_uri, ...)` is wrapped in `except Exception ...
logger.warning("[compile] could not mirror reel to %s: %s", ...)` ? the mirror failure is
swallowed, then the function unconditionally returns `{"status": "success", ...}`. The comment
right above it even admits: "On the L4 worker the mirror is how the reel reaches the bucket at
all (the worker's local file is discarded when the Job exits)" ? yet still swallows. The CP-
path justification ("the endpoint still falls back to the local file on the same instance")
does not hold on the worker, so the swallow is unguarded silent data loss there.
- `reel_uri` is non-empty on the worker: `cmd_compile` sets `config.storage.output_root` from
`--set/--out-uri` (worker/`__main__.py:485-486`), and `reel_object_uri` reads `output_root`. So
the mirror IS attempted and its failure IS reachable/swallowed.
- No verification anywhere downstream: `cmd_compile` (worker/`__main__.py:513-514`) sets `rc = 0
if status == "success" else 3` purely from the return value; `_write_compile_marker` copies the
status through; and control-plane `_dispatch_compile_remote` returns `{"status": "success",
download_url, ...}` whenever `result.returncode == 0`, with no `storage.exists(reel_uri)` check.
So a swallowed worker-side mirror failure propagates as a SUCCESS with a `download_url` the SPA
later polls, but the reel is neither in the bucket nor recoverable from the exited worker.

```

```

**Coverage gap is real (confirmed):** The real `_stitch_reel` (with its `storage.put` mirror +
swallow) is never executed by any test. All references mock it:
`tests/test_worker.py:1015,1062,1100` (all three `cmd_compile` tests) and
`tests/test_compile_dispatch.py:221` (CP in-process path). The dispatch/compute tests use
`_FakeCompileClient` (`tests/test_compile_dispatch.py:1108`) which "write[s] a reel done-marker
into the local bucket exactly where the real cmd_compile would" ? bypassing the real worker
entirely. No test injects a `storage.put` mirror failure to assert the resulting status/marker,
so the false-success-on-lost-reel path is entirely uncovered.

```

I could not refute this: both the failure mode and the absence of any covering test are concretely present in the merged code.

23. user (2026-07-08T06:18:04.580Z)

Structured output provided successfully