

claude-windows / 96963dab-d38f-488c-af44-530ddbfeed1

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-khelsutra-guru-badminton-highlight-indexer\96963dab-d38f-488c-af44-530ddbfeed1\subagents\workflows\wf_00bebd3c-82f\agent-a0f2bcb0972a20201.jsonl`
- SHA-256: `1ee9a680a157e6ced0cb381313f0107d131305669c6a385ff652dd3e3849572e`
- Source modified: `2026-07-08T05:51:52+00:00`
- Imported at: `2026-07-08T15:59:27+00:00`
- project: `wf_00bebd3c-82f`
- session_id: `96963dab-d38f-488c-af44-530ddbfeed1`
```

Transcript

1. user (2026-07-08T05:44:49.641Z)

You are mapping a subsystem of the repo at C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer (Windows checkout; read-only ? do NOT edit anything). Context: GitHub issue #524 wants to (a) move validation onto the L4 Cloud Run Job worker via a config-driven phase_placement knob, (b) drop a redundant re-hash in /api/process by trusting the upload's streamed video_id, and (c) design a topology where upload->stitch runs on the L4 with async copy-to-bucket. You are producing the precise code map a designer/implementer needs. Cite exact file paths and line numbers. Read the actual code, not just docs. Follow imports where the flow crosses files.

Subsystem: DOCS + CONVENTIONS. Read docs/COMPUTE_DECOUPLED_SERVING/ (the main design doc + runlogs dir listing), docs/PROJECT_DOC_TEMPLATE.md, docs/DOC_STATUS.md (lifecycle rules), and the ADRs for 012 (Cloud Run+GPU serving), 013 (GCP-only compute), 014 (compute-decoupled ratification) in docs/archives/decisions/DECISIONS.md. Summarize: (1) the structure/status-header conventions a NEW design doc under docs/COMPUTE_DECOUPLED_SERVING/ must follow, (2) the M0-M9 milestone framing of COMPUTE_DECOUPLED_SERVING and where a "phase placement / L4-primary" design fits, (3) any prior cost analysis (Cloud Run GPU pricing, per-run cost numbers) already in the docs that the new doc should reuse or update, (4) exact constraints from ADR-012/013/014 that bound the topology options (e.g. GCP-only, Cloud Run chosen over GCE ? why, and is that ADR binding for serving topology?).

2. user (2026-07-08T05:45:02.508Z)

```
architecture.svg
README.md
runlogs
TESTING_STRATEGY.md
```

3. user (2026-07-08T05:45:09.630Z)

```
1      # Compute-Decoupled Rally Serving ? one pure core, profile-selected compute + storage
2
3      > **Status:** ? **M0 APPROVED** (owner, 2026-06-21) ? **IN PROGRESS (build M1+)** .
**Owner:** Avi . **Created:** 2026-06-21 . **Last updated:** 2026-06-21
4      > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to
`docs/archives/past_projects/` when DONE)
5      > **Tracking anchors:** $7 progress tracker is the source of truth; pointer in memory
`current-state` + `docs/NEXT_STEPS.md` once work starts.
6      > **Relation to existing docs:** the productionization successor to the archived
`DECOUPLED_COMPUTE` (M3/M4); realizes the `ComputeBackend`/`StorageBackend` registries from
[PLATFORM_ARCHITECTURE.md](../PLATFORM_ARCHITECTURE.md) $3; the **truly-local** profile is the
L0 tier of [VIDEO_LOCALITY_MODEL.md](../VIDEO_LOCALITY_MODEL.md) and the `LocalDesktop` backend
of [DESKTOP_APP_PLAN.md](../DESKTOP_APP_PLAN.md).
7      > **Naming:** this was the "Cloud Run + GPU serving subproject" (ADR-012). Renamed
**compute-decoupled serving** because **truly-local is a co-equal, first-class profile**, not an
afterthought.
```

```

8      > **Honesty note:** claims tagged `[verified]` (checked against code by the design
workflow `wf_56274ff0`), `[design]` (proposed), `[researched]`.
9
10     ---
11
12     > **? Northstar (D14):** a **mobile, app-like** flow ? point at a **Google Drive** video
? **fully cloud-native** run ? the stitched, **ready-to-share** reel lands **back in Drive, next
to the source**. Two design promises around it: a user can **switch local?cloud anytime** (D13,
a headline feature), and execution is **never a surprise** ? auto-detect *suggests*, the user
*confirms* (D15).
13
14     ## 0. TL;DR
15     The rally engine is a **pure function of `(input bytes + config) ? outputs`** in three
deterministic stages ? **DETECT ? WINDOW ? STITCH** ? that **no deployment profile may branch
on**. A **`ComputeBackend` (deployment profile)** decides *where* the core runs (a local process
vs a Cloud Run GPU job); a **`StorageBackend`** decides *where bytes come from* (local FS / USB
/ mounted Google Drive / bucket). **One startup config** picks both. We ship **truly-local**
(local GPU + in-process stitch, zero cloud) and **cloud-serving** (upload `<2GB` / gdrive /
bucket ? a Cloud Run GPU job runs detect **and** stitch, so stitch's compute is metered), plus
**cpu-mock** for CI. The single most important caveat: detection is **byte-identical only on the
same GPU** (cross-GPU is sub-pixel-different, per the 2026-06-20 deploy report) ? so the
**parity guarantee is enforced at the WINDOW + STITCH layer**, not detector-CSV bytes.
16
17     ## 1. Problem & goal
18     - **Why now:** DECOUPLED_COMPUTE proved compute *portability* (M0?M2 + M3.5 on a GCP L4)
and is archived; its deferred **M3 (serving endpoint + per-video billing) + M4 (scale-to-zero +
cost guard)** become this project. The owner needs **both** a hosted service *and* a truly-local
mode, with the **core (detect+stitch) unaffected** by which one runs.
19     - **The two user-facing outcomes** (owner, 2026-06-21):
20     - ** (a) Hosted service** ? operate on small **local uploads (`<2GB`)** *and* load from
**gdrive / a bucket**, spinning up **Cloud Run for both rally segmentation AND stitching**
(stitching is compute-intensive ? run it where its cost is accounted for).
21     - ** (b) Truly-local** ? on a local machine with videos on **local drives/USB + a local
GPU**, run inference **and** stitching **completely locally**.
22     - **"Good" looks like:** the same `video ? rallies ? reel` core gives identical windows
+ stitched ranges whether it runs on a laptop or Cloud Run; switching is a **config/startup
choice**, never a code branch in the core.
23
24     ## 2. Decisions locked
25
26     | # | Decision | Source | Implication |
27     |---|-----|-----|-----|
28     | D1 | The core is a **pure 3-stage function** (DETECT/WINDOW/STITCH); **no profile
branch inside it**. | design `[verified]` | All profile difference lives in config + the two
registries. |
29     | D2 | **Two registries:** `ComputeBackend` (where) selected by `deployment.profile` +
`compute_target`; `StorageBackend` (what-bytes) by `input_backend`/`storage.backend`. | ADR-014
| Realizes PLATFORM_ARCHITECTURE's `local`/`hosted`/`byo_worker`. |
30     | D3 | **Profiles shipped:** `truly-local`, `cloud-serving`, `cpu-mock` (CI);
`byo_worker` reserved/deferred. | ADR-014 | Enum/seam reserved so BYO isn't precluded. |
31     | D4 | In **cloud-serving, STITCH runs on Cloud Run** (co-located with detect) so
CPU/wall + egress are metered into the per-job cost ledger. | owner (a) | Stitch-on-laptop would
hide real cost. |
32     | D5 | **Parity enforced at WINDOW+STITCH** (byte-exact for a fixed trajectory);
**detect** is byte-exact only same-GPU. | deploy report 2026-06-20 | Cross-GPU asserted at the
rally/window level, never CSV bytes. |
33     | D6 | **Default-OFF, smallest-safe-first, one PR per milestone**; any core-caller
change (e.g. configurable output base) ships behind a flag + a Launch Report. | [[launch-report-
convention]] | Zero-regression discipline. |
34     | D7 | **Cloud Run cold-start: SCALE-TO-ZERO** ? accept the ~1?3 min cold-start (it's an
async job: upload?poll?reel, amortized over the multi-minute job). **No warm pool** (that
~$500/mo idle would break per-video billing). | owner 2026-06-21 | M6/M7. Revisit a warm lane
only on a UX complaint. |
35     | D8 | **Pricebook = a versioned DB table**; cost computed in **USD** (matches GCP
billing = one source of truth), **displayed in INR** at a configurable FX rate (Bangalore

```

market); re-versioned when GCP prices change. | owner 2026-06-21 | M8. |

36 | D9 | **Edge auth = Cloudflare Access** (reuse the site's existing CF). Lock the Cloud Run **ingress** to the CF proxy + **verify the `Cf-Access` JWT signature** (so a direct hit to the Cloud Run URL can't spoof the identity header). | owner 2026-06-21 | M9. One identity system. |

37 | D10 | **Free-tier "data-for-reels" trade:** free reels in exchange for **training rights** (uploaded footage + derived trajectories/labels); **paid tier = no-train** (privacy is the paid feature). **Carve-outs:** train ONLY on attested **ADULT** footage (minor footage ? reel yes, training pool **NO** ? DPDP/GDPR need verifiable parental consent); train on **DERIVED** data + **auto-delete raw**; **ToS counsel-reviewed**; live training-on-uploads **gated to M9** (public signup). Truly-local needs no DPA. | owner 2026-06-21 | M9 + D12; ties [[paper-coauthor-aim]] / PERSONALIZATION flywheel. |

38 | D11 | **Quality floor: Gemini handoff ON** for cloud-serving until the owned model clears the #172 ship-gate (e.g. ?0.70 multi-court) ? then flip via a **Launch Report**; the owned model runs in **shadow/eval** meanwhile. (Gemini = serving/inference ONLY, never a training source ? CLAUDE.md guardrail.) | owner 2026-06-21 | M7 + D12. |

39 | D12 | **Data-flywheel harness is IN SCOPE** (owner add, Q5): **capture the (consented, adult) uploaded video + the Gemini reel/rally output** ? **reliably store** in the per-video workspace/bucket ? **run shadow evals** (owned-vs-Gemini, dual-eval-set) ? **promote good ones to the golden TRAINING set** (human-reviewed labels) so the owned model climbs toward the D11 floor (then Gemini flips off). | owner 2026-06-21 | new **M11**; ties D10 (consent) + `data\_pipeline/` + the eval/experiment harness. |

40 | D13 | **Profile is user-switchable ANYTIME** ? a headline **FEATURE**. The same user runs **local today, cloud tomorrow, back to local** ? it's just a config/startup choice; the pure core gives **equivalent** output either way. **Advertise it** ("your machine *or* our cloud ? your call, switch anytime"). | owner 2026-06-21 | Honest caveat: equivalent at the **rally/window** level, not bit-identical across GPUs (D5); a single job runs in one profile, switching is **between** jobs. |

41 | D14 | **NORTHSTAR** ? operate toward this: **a mobile, app-like** flow ? point at a **Google Drive** video ? **fully cloud-native** run ? the stitched, **ready-to-share** reel lands **back in Drive**, next to the source. | owner 2026-06-21 | Implies a **gdrive-api** (OAuth) **StorageBackend** (read source + write reel back) ? distinct from the local **mount** backend; **resolves \$8 #7**, new **M12**. The mobile UI is the khelsutra track; the engine must support it (async jobs + status + Drive round-trip / signed URL). Net-new scope (OAuth + Drive write); not necessarily v1, but the design must not preclude it (the **StorageBackend** ABC accommodates it). |

42 | D15 | **Auto-detect SUGGESTS**, the user **CONFIRMS** ? execution is **NEVER** a surprise. A GPU owner may still choose cloud; **cloud (= spend)** is never entered silently; the active profile is always explicit + surfaced. | owner 2026-06-21 | Amends \$4.3 (auto-detect = a proposed default, not an auto-decision). |

43 | D16 | **M6/M7 implementation SHAPE** + the cheap \$8 picks (owner 2026-06-21): (a) **Cloud Run JOBS** (not Services-with-GPU) for the detect+stitch unit, **bucket-poll** for completion (reuses `JobWaiting`/`claim_due_job`` verbatim, fits async upload?poll?reel + scale-to-zero D7); (b) **stitch ENCODE = libx264** (protects the D5 byte-determinism/parity guarantee + the L4 parity test) ? **NVDEC for DECODE only** (a speed lever that never touches output bytes); **NVENC encode is DEFERRED** behind a config knob, revisit only if M7 cost data justifies relaxing parity (\$8 Q8); (c) **input cap = hard-reject <2GB** for v1 (a config cap so raising to 4GB later is one line; matches the shipped M2 chunked-upload threshold ? \$8 Q9). | owner 2026-06-21 | M6/M7. Resolves \$8 Q2 (co-located, = D4), Q8, Q9, Q10 (dir name confirmed). |

44 | D17 | **Real cloud verification BUDGET = \$100 / ?10k cap** (owner 2026-06-21): the owner authorized actual GCP spend for due-diligence + production-grade testing/verification of M5?M7 (real L4 run, the M5 runlog, the M7 DEPLOY_REPORT), **provided the cap is respected** ? confirm exact cmd + hourly cost before each spend, prefer the smallest viable run + cheapest-adequate GPU, and **DELETE every resource after** ([[gcp-vm-always-delete]]). Removes the M6/M7 "spend" gate within the envelope; the calibrated pricebook (M8) + counsel ToS (M9) + OAuth app (M12) stay owner/counsel-side. | owner 2026-06-21 | Unblocks the real M5/M7 runs (not just the code) within the cap. |

45

46 **## 3. Foundation ? evolution / ADR lineage** ? **trace the idea end-to-end**

47 This project is the latest step in a long decision chain; each ADR contributed a piece and a **subtlety that carries forward**:

48

49 | ADR / doc | Contributed | Subtlety carried into this project |

50 |---|---|---|

51 | **ADR-005** | Permissive licensing (MIT/Apache-2.0/BSD only) | The engine + any served

model + the **container image** (ffmpeg, fonts, weights) stay commercial-clean. |

52 | **ADR-008** | Owned-model topology; **train==serve** featurizer single-sourcing | The "core unaffected across profiles" is the same parity discipline, extended from train/serve to **local/cloud**. |

53 | **ADR-009** | KhelSutra Guru; **local-first delivery** | The **truly-local profile** is the direct continuation ? privacy/offline self-host stays first-class. |

54 | **ADR-010** | DECOUPLED: D0?GCP pivot | The cloud-compute origin (GPU + co-located bucket). |

55 | **ADR-011** | DECOUPLED scope = M0?M2+M3.5; **deferred M3 (serving) + M4 (billing/cost-guard)**; overrode the "rent nothing / not a service" posture | **This project realizes M3/M4.** Carries \$06's owner-gated gates: **DPA/consent** for non-owner uploads, collector `md5` keying, WASB-C3. |

56 | **ADR-012** | GPU-native; TPU deferred; **Cloud Run+GPU scale-to-zero = serving model**; **NVDEC** named as the decode lever | The serving substrate + the 7-step seed scope this expands. |

57 | **ADR-013** | Drop D0; **GCP sole compute stack**; \$200 D0 credits ? non-compute only | The provider for cloud-serving; D0 Spaces remains only a possible S3 **storage** target. |

58 | **ADR-014 (new)** | **Compute-decoupled serving**: one pure core, profile-selected compute+storage; stitch-where-metered | This doc. **Refines** ADR-011/012, does **not** supersede them. |

59 | PLATFORM_ARCHITECTURE.md | The `ComputeBackend`/`StorageBackend` registries | The abstraction this realizes in code. |

60 | VIDEO_LOCALITY_MODEL.md / DESKTOP_APP_PLAN.md | L0 fully-local tier / `LocalDesktop` backend | The truly-local profile = these, productionized. |

61 | 07-COMPUTE-ABSTRACTION (archived) | `DeviceContext`/`DeviceProfile` (designed) | WASB hardcodes `device='cuda'` with **no CPU fallback** ? a portability gap M4 closes. |

62 | DEPLOY_REPORT_GCP_2026-06-20 (archived) | First real L4 run; **cross-GPU sub-pixel parity finding** | Why D5 (parity at window level, not CSV bytes). |

63

64 **## 4. Design / architecture**

65 See architecture.svg for the picture. **The core never changes; the box around it does.**

66

67 **### 4.1 The core contract (must stay byte/behaviour-identical) [verified]**

68 1. **DETECT (GPU)**: `DetectRunner.run_predict(video_path, output_dir, video_id) ? CSV[Frame,Visibility,X,Y]` (`backend/pipeline/detectors/base.py:164`). Three interchangeable impls ? `WasbRunner` (WSL), `NativeWasbRunner` (Linux GPU), `StubDetectorRunner` (CPU mock) ? emit the **identical CSV schema**. The trajectory CSV is the stable artifact boundary.

69 2. **WINDOW (CPU pure fn)**: `trajectory_to_action_windows(points, fps, frame_width, ?) ? windows` (`tracknet\_runner.py:281`). A shared `@staticmethod`, zero GPU/IO/randomness ? same trajectory + config = byte-identical candidate windows. Used by every runner **and** the eval/regression stack verbatim.

70 3. **STITCH (CPU/ffmpeg)**: `VideoCompiler.compile_highlights(raw_video, segments, out_name) ? reel` (`stitcher/compiler.py:303`). Pure FS+subprocess (merge ? ffmpeg ? watermark ? per-clip libx264 trim ? concat). Deterministic for a given input + ranges + watermark config.

71

72 **Non-goal (the trap to avoid)**: **no** code branch inside detect/window/stitch keyed on profile ? the same discipline the experiment harness already enforces (a disabled cue is byte-identical to CONTROL).

73

74 **### 4.2 The deployment profiles**

75 | Profile | Compute | Storage | Orchestration | When |

76 |---|---|---|---|---|

77 | **truly-local** | `compute_target=local_gpu`; `detector_impl=native` (Linux) / `auto` (WSL); stitch **in-process** | `local` (or `gdrive` = mounted Drive) | `python -m backend.main` or `worker infer` CLI; existing async job worker | Self-host / offline / privacy / dev-on-a-GPU-box ? owner (b) |

78 | **cloud-serving** | `compute_target=cloud_run`; `detector_impl=native` (L4, cuda); **detect AND stitch** on the same Cloud Run job | `gcs` (or `gdrive` mount / assembled upload); per-video workspace; reels via signed URL | Cloud Run control plane enqueues ? Cloud Run GPU job pulls/runs/pushes; `WAIT_AND_RESUME` reschedules the control-plane job; scale-to-zero | Hosted SaaS ? owner (a) |

79 | **cpu-mock** | `compute_target=cpu_mock`; `detector_impl=stub` (synth/replay) | `local/in-memory` fakes | `run_all_tests.py` / pytest | CI, regression, profile-parity, GPU-free dev |

```

80      | *byo_worker* (deferred) | tenant GPU via outbound pull agent | tenant's own bucket
(signed URLs) | long-poll pull; same job table scoped by tenant | Privacy-sensitive enterprise ?
**DEFERRED**, enum/seam reserved |
81
82      ### 4.3 Profile selection (the "startup/setup config") `[design]`
83      ONE new top-level `deployment` section on `AppConfig` + preset application + env auto-
detect in `load_config`:
84      - `deployment.profile` ? {truly-local | cloud-serving | cpu-mock}` (empty = today's
manual knobs, **fully backward-compatible**).
85      - Selecting a profile applies a **coherent preset bundle** of *existing* knobs that
today drift independently: **INPUT** (new `input_backend`/`input_root`, parallel to the output-
side `StorageConfig`), **COMPUTE** (new `compute_target` enum locking `detector_impl` +
`skip_ai_handoff` + workspace strategy), **STORAGE**
(`storage.backend`/`workspace_root`/`single_folder_per_video`, all already exist).
86      - **Precedence:** built-in defaults ? profile preset ? `config.json` ? env
(`RALLY_DETECTOR_IMPL`, `WASB_*`, `DEPLOYMENT_PROFILE`, ?) ? worker `--set k=v`. Every knob
stays individually overridable.
87      - **Env auto-detect SUGGESTS, the user CONFIRMS (D15 ? no surprise execution):** the env
gives a *proposed* default (`K_SERVICE` ? cloud-serving; `CI=true` ? cpu-mock; else
`torch.cuda.is_available()` ? truly-local), but the user **explicitly confirms or overrides** it
? a GPU owner may still pick cloud, and **cloud (= spend) is NEVER entered silently**. The
active profile is always surfaced (logged/shown). **Per-profile startup checks fail/warn fast**
(cloud-serving without GCS creds/GPU ? loud error before any job).
88      - The worker (`cmd_infer`/`cmd_train`) is **already profile-agnostic** (reads
`config.storage`, accepts `--config`/`--set`) ? no worker rewrite.
89
90      ### 4.4 Compute placement ? and *why stitch runs on Cloud Run* `[design]`
91      DETECT on GPU everywhere; WINDOW CPU-pure everywhere; STITCH is real CPU work. In
**cloud-serving, stitch is co-located in the Cloud Run job** so `gpu_active_seconds` (detect) +
`wall_seconds` (detect+stitch) + `bytes_out` (reel egress) land on **one metered invocation** ?
the per-job cost ledger. Running stitch on a laptop would leave that cost
**invisible/unattributed** ? exactly what the owner wants to avoid. (Stitch stays synchronous in
the GPU job for the common one-reel case; a queued CPU-only stitch fallback only if compiles get
bursty ? the job table + `claim_due_job` already supports both.)
92
93      ### 4.5 Reuse map ? most of the seams already exist `[verified]`
94      **Exists:** the 3-stage core; `_resolve_runner` (the single compute seam);
`StubDetectorRunner` (`replay_csv` = byte-exact $0 anchor); `StorageBackend` ABC + 4 backends +
`StorageRef.parse_uri`; the worker CLI (fetch/put, `--config`/`--set`); the async job control
plane (`jobs` table, `claim_due_job` CAS); `ExecutionPolicy WAIT_AND_RESUME`/`JobWaiting`
(carries to Cloud Run dispatch verbatim); the chunked-upload router (<2GB); `skip_ai_handoff`
(LocalNoAI); storage fakes; golden fixtures + experiment harness. **Partial:** per-video
workspace helper (default-OFF); `DeviceContext` (designed); edge-auth (v4 `user_id` columns
exist). **Net-new:** the `deployment` section; `input_backend`; configurable output base (kill
hardcoded `output`); Cloud Run dispatch shim + container image; cost ledger + pricebook; edge-
auth header extraction; startup env detection.
95
96      ### 4.6 The data-flywheel harness (D12 / M11) `[design]`
97      Gemini-on (D11) ships good reels now; this harness is **how we earn the right to flip
Gemini off**. On a **consented, adult** cloud-serving job (the D10 trade), the pipeline
additionally:
98      1. **Captures** the source upload + the Gemini reel + the rally output
(intervals/report) into the per-video workspace (`workspaces/<video_id>/`), instead of the no-
consent auto-delete path.
99      2. **Reliably stores** them durably in the bucket (the consented-retention bucket; raw
video kept only for the training pool, not the no-consent path).
100     3. **Runs shadow evals** ? the owned model scores the same video alongside Gemini
(`backend/eval/experiment.compare` + the dual-eval-set, [[eval-dual-set-regression]]) ? tracks
the owned-vs-Gemini gap per job (the live read on "is the floor reachable yet?").
101     4. **Promotes to golden** ? captured videos that pass review become **golden TRAINING
rows** (human-reviewed labels via the annotation flow ? `data_pipeline/`), growing the corpus ?
the owned model climbs to the D11 floor ? flip Gemini off via a Launch Report.
102     This is the consent?capture?eval?goldenize **flywheel** (the cloud realization of the
PERSONALIZATION contribution model): every consented free reel makes the owned model better.
Reuses `data_pipeline/`, `backend/eval/`, and the M9 consent gate; **adults-only**, raw-

```

retention gated by the D10 consent.

103

104 ## 6. Honest limits ? what this does NOT do / prove

105 - A green CI suite proves ****plumbing + determinism (window+stitch) + profile-parity****
for \$0 ? it does ****NOT**** prove field quality, nor that real-GPU detection is good (that's owner-
side, real-video).

106 - ****Cross-GPU detection is NOT byte-identical**** ? parity is a window-level claim, not a
CSV-byte claim.

107 - The cloud `0.1.0-l4` weights are ****n=2 plumbing-grade**** ? cloud-serving ships with
****Gemini handoff ON**** until the owned model clears a floor (open Q).

108 - This does not resolve ****WASB-C3**** (sharing a trained model) or the ****DPA/consent****
gate ? both remain owner/counsel-gated before public, non-owner serving.

109

110 ## 7. Deliverables & progress tracker ? ****source of truth****

111 Legend: ? Todo . ? In progress . ? Done . ? Gated. ****One small PR per row****, off
`origin/master`, no stacking. Default-OFF where it touches core callers.

112

113 | ID | Deliverable | Depends | Gated? | Status | PR |

114 |----|-----|-----|-----|-----|----|

115 | M0 | This design doc + project dir + `runlogs/` + ADR-014 | ? | owner sign-off | ? |

****? APPROVED 2026-06-21**** |

116 | M1 | `deployment.profile` + `compute\_target` config + preset bundles +
precedence/auto-detect; unit tests (no behaviour change, profile='` default) | M0 | safe | ? |
[#279](https://github.com/Khelsutra/badminton-highlight-indexer/pull/279) |

117 | M2 | `input\_backend`/`input\_root`; wire main CLI/API input through
`StorageRef.parse\_uri` (local=identity); storage-fake tests | M1 | safe | ? |

[#280](https://github.com/Khelsutra/badminton-highlight-indexer/pull/280) |

118 | M3 | Configurable output/scratch base ? kill hardcoded `output/`
(`trajectory\_hybrid:237,313`, `yolo\_hybrid:407`, `gemini`); ****DEFAULT** stays `output/`; golden
+ stub-E2E byte-identical | M1 | ? default-OFF + Launch Report on flip | ? |
[#282](https://github.com/Khelsutra/badminton-highlight-indexer/pull/282) |

119 | M4 | `DeviceContext` + WASB ****CPU-fallback****; unified healthcheck/device wiring | M1 |
safe (cuda default unchanged) | ? | [#283](https://github.com/Khelsutra/badminton-highlight-
indexer/pull/283) |

120 | M5 | ****truly-local profile**** end-to-end (preset + startup checks); first committed
****runlog**** (truly-local sample run) | M2,M3,M4 | owner real-GPU | ? |

[#286](https://github.com/Khelsutra/badminton-highlight-indexer/pull/286) (preset+startup-
checks+L4 parity ?; runlog ? owner GPU) |

121 | M6 | Cloud Run GPU ****dispatch shim**** (control plane ? job via storage ? re-claim) +
****containerized worker image**** (ffmpeg+CUDA+WASB+fonts); mocked-dispatch tests | M5 | ? owner
(GCP spend) | ? | [#287](https://github.com/Khelsutra/badminton-highlight-indexer/pull/287)
(shim+image+mocked ?; real build/run = M7 ? owner) |

122 | M7 | ****cloud-serving profile**** e2e on L4: upload`<2GB` + gdrive/bucket ? detect+stitch
on Cloud Run ? reel via signed URL; ****DEPLOY_REPORT**** (posterity, sample runs + contact sheet) |
M6 | ? owner (spend) | ? | ? |

123 | M8 | Per-job ****cost ledger**** (v6 migration: `owner_id, gpu_active_seconds,
wall_seconds, bytes_out, cost_usd, cost_breakdown_json`) + ****pricebook**** + aggregation +
`/api/.../cost` | M7 | ? owner (billing) | ? | [#288](https://github.com/Khelsutra/badminton-
highlight-indexer/pull/288) (ledger+pricebook+/cost ? default-OFF; real prices + usage-emission
? owner/follow-up) |

124 | M9 | ****Edge auth**** (Cf-Access extraction) + per-tenant scoping on `user\_id`;
DPA/consent gate wiring | M7 | ? owner (privacy/counsel) | ? |

[#290](https://github.com/Khelsutra/badminton-highlight-indexer/pull/290) (auth: Cf-Access JWT
verify + owner_id tag, default-OFF ?; ****+ exposure-safety boot posture**** in `startup.py` ? edge-
auth-on requires the `allowed\_video\_root` jail + CF team/AUD or the server refuses to start, see
changelog; consent cols + ToS/DPA ? owner/counsel) |

125 | M10 | `byo\_worker` / zero-infra-BYO ? ****design-only, DEFERRED**** | M8,M9 | ? explicit
owner approval | ? | ? |

126 | ****M11**** | ****Data-flywheel harness (D12):**** on a consented adult job, ****capture**** the
upload + Gemini reel/rally output ? ****reliably store**** under the per-video workspace ? ****shadow-
eval**** owned-vs-Gemini (dual-eval-set / `experiment.compare`) ? ****promote**** good ones to the
golden TRAINING set (human-reviewed labels via the annotation flow). Closes the loop so the
owned model climbs to the D11 floor. Reuses `data\_pipeline/` + `backend/eval/` + the consent
gate (M9). | M7,M9 | ? owner (consent/privacy) | ? | ****skeleton landed**** (see changelog):
`backend/flywheel.py` + `FlywheelConfig` (default-OFF) + gated `\_maybe\_run\_flywheel` hook,

```

**consent faked-deny**; live activation ? owner/counsel (consent schema) + owned-model/golden
data |
127 | **M12** | **`gdrive-api` (OAuth) StorageBackend (D14 Northstar):** cloud reads the
source from + writes the stitched reel **back to the user's Google Drive next to the source**
(per-user OAuth consent) ? the mobile/cloud-native Drive round-trip. Distinct from the local
mount backend; slots into the `StorageBackend` ABC. | M7,M9 | ? owner (OAuth/consent) | ? |
**backend landed** (see changelog): `backend/storage/gdrive_api.py` `GDriveApiStorage` +
`gdrive-api://` scheme + ADC build; hermetic fake-Drive tests. **Live per-user OAuth app ?
owner** |
128
129 **Testing strategy is a first-class deliverable** ?
[ `TESTING_STRATEGY.md` ](TESTING_STRATEGY.md). **Run logs / sample runs for posterity** ?
[ `runlogs/` ](runlogs/).
130
131 ## 8. Open questions ? owner *(blocking-gates vs nice-to-knows)*
132 **? The 5 gating questions are LOCKED (owner, 2026-06-21) ? see §2 D7?D12:**
133 1. ? **Cold-start ? D7** (scale-to-zero, accept cold-start).
134 3. ? **Pricebook ? D8** (versioned DB table; USD internal / INR display).
135 4. ? **Edge auth ? D9** (Cloudflare Access; verify Cf-Access JWT, lock ingress to
proxy).
136 5. ? **DPA/consent ? D10** (free-tier data-for-reels trade; adults-only; derived-data;
counsel-gated to M9).
137 6. ? **Quality floor ? D11** + the **data-flywheel harness D12 / M11** (Gemini-ON until
the owned model clears the floor; capture?store?eval?goldenize meanwhile).
138
139 **? The remaining smaller picks are LOCKED (owner, 2026-06-21) ? see §2 D16:**
140 2. ? **Stitch placement ? co-located** in the detect job (= D4; one metered unit; the
job table supports a future CPU-split non-breaking, revisit only on bursty-compile metrics).
141 7. ? **RESOLVED ? D14:** cloud needs a **`gdrive-api` (OAuth) StorageBackend** (read the
source from + write the reel back to the user's Drive) ? the Northstar's Drive round-trip ?
distinct from the local **mount** backend (which stays for truly-local). New milestone **M12**.
142 8. ? **Codec ? D16:** **libx264 ENCODE** (protect parity) + **NVDEC DECODE-only**
(speed, no output-byte change); NVENC-encode deferred behind a config knob.
143 9. ? **Input cap ? D16:** **hard-reject`<2GB`** for v1, as a config cap (raise to 4GB
later in one line).
144 10. ? **Directory name ? `COMPUTE_DECOUPLED_SERVING/` CONFIRMED final** (already the
name everywhere + in ADR-014); `CLOUD_RUN_GPU_SERVING/` dropped.
145
146 ## 9. Definition of done
147 - [ ] M1?M4 land (config + input + output-base + device layer), each green on golden
cross-OS + experiment no-regression.
148 - [ ] **Profile-parity test green** ? truly-local vs cloud-mock produce byte-identical
windows + segment ranges.
149 - [ ] **truly-local** (M5) and **cloud-serving** (M7) each have a committed **runlog /
DEPLOY_REPORT** with a real sample run.
150 - [ ] Cost ledger (M8) makes GPU + **stitch** + egress cost visible per job; edge auth
(M9) gates multi-tenant.
151 - [ ] **Data-flywheel harness (M11)** captures?stores?shadow-evals?goldenizes consented
adult jobs; the owned model is on a measured path to the D11 floor.
152 - [ ] ? flip Status DONE, archive to `docs/archives/past_projects/`.
153
154 ## 10. References
155 **ADR lineage:** ADR-005/008/009/010/011/012/013 ? **ADR-014**
(`docs/archives/decisions/DECISIONS.md`). **Strategy:** `PLATFORM_ARCHITECTURE.md`,
`HOSTING_PLAN.md`, `VIDEO_LOCALITY_MODEL.md`, `DESKTOP_APP_PLAN.md`,
`LOCAL_APPLIANCE_ARCHITECTURE.md` (kheIsutra). **Archived predecessor:**
`docs/archives/past_projects/DECOUPLED_COMPUTE/` (esp. `05-COST-ATTRIBUTION.md`, `07-COMPUTE-
ABSTRACTION.md`, `06-RISKS-AND-OPEN-QUESTIONS.md`, `DEPLOY_REPORT_GCP_2026-06-20.md`). **Code
anchors`[verified]`:** `backend/pipeline/detectors/base.py:164`, `tracknet_runner.py:281`,
`stitcher/compiler.py:303`, `segmenters/trajectory_hybrid.py:65,237`, `backend/storage/*`,
`backend/worker/__main__.py:88`, `backend/api/{uploads,job_worker}.py`,
`backend/infrastructure/database.py`. **Design provenance:** workflow `wf_56274ff0` (4 code-
mappers ? synthesis; 58 findings).
156
157 ---

```

```

158     ### Changelog
159     - 2026-06-21 ? created from the design workflow output; M0 in progress.
160     - 2026-06-21 ? **M0 APPROVED**; **M1 landed**
([#279](https://github.com/Khelsutra/badminton-highlight-indexer/pull/279)):
`deployment.profile` + `compute_target` + preset bundles + precedence/auto-detect-as-suggestion
(D15) on `AppConfig`. Pure/additive, `profile=''` default == today's behaviour (zero core
change); presets touch only existing knobs (no M2 `input_backend` / M3 output-base over-reach).
Suite 1068 green, cov 84.11%.
161     - 2026-06-21 ? **M2 landed** ([#280](https://github.com/Khelsutra/badminton-highlight-
indexer/pull/280)): `input_backend`/`input_root` + URI-aware main API/CLI input via
`resolve_input_uri`/`acquire_local_input` (local = identity passthrough, byte-for-byte
unchanged; bucket/Drive fetched to scratch). Unsupported scheme ? clean 400; `local://` accepted
on the resolved key; hosted-mode rejects non-local URIs; source URI kept for DB/report
provenance. Suite 1105 green, cov 84.84%.
162     - 2026-06-21 ? **M3 landed** ([#282](https://github.com/Khelsutra/badminton-highlight-
indexer/pull/282)): configurable output/scratch base via
`backend/utils/paths.output_base(config)` ? the 3 segmenters' chunk/handoff/temp dirs root under
`output.output_dir` (default `"output"`, byte-identical; default-OFF ? no Launch Report). Fixes
the latent bug where scratch ignored `--output-dir`. Suite 1114 green, cov ~84.9%.
163     - 2026-06-21 ? **M4 landed** ([#283](https://github.com/Khelsutra/badminton-highlight-
indexer/pull/283)): `DeviceContext` + native WASB **CPU-fallback** via a new `device="auto"`
(cuda-if-available-else-cpu); explicit `"cuda"` stays strict (no silent slow-CPU downgrade),
default unchanged. `"auto"` resolves before argv (never reaches `wasb_infer.py`); CUDA probe
cached. Device-selection tested offline; real CPU-inference run owner/GPU-verified. Suite 1126
green, cov 84.96%. **? M1?M4 (the config/input/output/device layer) COMPLETE.**
164     - 2026-06-21 ? **M5+ batch authorized + design picks locked** (owner): the remaining $8
picks resolved ? **D16** (M6 = Cloud Run **Jobs** + bucket-poll; **libx264 encode + NVDEC
decode-only**; **hard-reject <2GB** input cap; dir name confirmed) and **D17** (real cloud
verification budget = **$100/?10k cap**, delete-after). Owner gave a blanket go to write+merge
the **default-OFF code** for every remaining gated milestone (M5,M6,M8,M9-auth,M9-consent-
cols,M11-plumbing,M12), one isolated PR each, reserving the gate for the real
run/spend/calibration/counsel. Grounded by gate-map workflow `wf_4849400f`.
165     - 2026-06-21 ? **REAL M7 cloud GPU run ?** (owner-authorized, within the D17 cap):
provisioned an **L4** (`g2-standard-4`, asia-southeast1-a; Mumbai stocked-out) ? native WASB on
the L4 ? **22 rallies** on the smoke clip ?
`gs://?/outputs/05e0e577?/{intervals.csv,report.json}`; VM DELETED ($0, ~$0.21). ? the no-Gemini
cloud run needs `--set indexing.skip_ai_handoff=true` (the fully-local path emits WASB candidate
windows directly ? the long-flagged "0 rallies" cold-start fix). The reel **stitch** validated
locally (`VideoCompiler` ? 22-segment highlights, 85 MB). DEPLOY_REPORT runlog = the cloud-GPU
dry-run guide (`docs/CLOUD_GPU_DRYRUN_GUIDE.md`).
166     - 2026-06-21 ? **M9-auth code landed** ([#290](https://github.com/Khelsutra/badminton-
highlight-indexer/pull/290)): `backend/api/auth.py` ? Cloudflare Access **Cf-Access JWT verify**
(RS256: sig+exp+iss+aud against the team JWKS; lazy `PyJWT[crypto]`; injectable JWKS ? fully-
offline tests) + `resolve_tenant` ? tags `create_job(owner_id=?)`;
`SecurityConfig.{edge_auth_enabled=False,cf_team_domain,cf_access_aud}` + CORS via
`RALLY_CORS_ALLOW_ORIGINS`. **Default-OFF.** Security review (`wf_9b4c71d0`) found NO bypass;
locked alg=none + HS256-confusion rejection. **? M9 ? ? the consent cols + ToS/DPA are
owner/counsel; the consent migration is now v7 (M8 took v6).**
167     - 2026-06-21 ? **M8 code landed** ([#288](https://github.com/Khelsutra/badminton-
highlight-indexer/pull/288)): per-job **cost ledger** ? `backend/billing.py` (pure `usage x
pricebook` math, USD internal / INR display) + a **v6 migration** (NULLABLE cost columns on
`jobs` + a versioned `pricebook` table seeded `placeholder-v0` at $0; additive-NONE, idempotent)
+ `record_job_cost`/`get_job_cost`/`get_video_cost` + `GET /api/jobs/{id}/cost` +
`/api/videos/{id}/cost` + `BillingConfig` (**default-OFF**) + a gated `_maybe_record_job_cost`
hook. Adversarial review (`wf_498b24af`) folded a **major silent-GPU-under-bill** (an unpriced
`gpu_type` is now surfaced, not invisible $0) + the **honest no-op** caveat (usage emission is a
follow-up ? cost 0 until wired) + breakdown reconciliation. Suite green, ruff+mypy clean. **? M8
? ? owner inserts real GCP prices (a new pricebook version) + the worker usage-emission is a
follow-up; placeholder ships $0 = no user-facing change. v6 is the shared bump (M9-consent cols
extend it).**
168     - 2026-06-21 ? **M6 code landed** ([#287](https://github.com/Khelsutra/badminton-
highlight-indexer/pull/287)): `backend/compute/` ComputeBackend registry ? `CloudRunCompute`
dispatches a Cloud Run **Job** (injectable client) then **bucket-polls** a done-marker via
`execution_policy.poll_until` + the storage layer (D16); `get_compute_backend` returns a backend
ONLY for `compute_target=cloud_run` (else `None` ? the worker's in-process path, **default-

```


OFF**). The dispatched container runs **inline** (forced `local_gpu+native`, never re-dispatch); the worker writes a `--done-uri` completion marker on success.

`deploy/cloudrun/{Dockerfile,README}` = the L4 image (CUDA+ffmpeg+fonts+wasb` conda env; weights staged per WASB-C3) + the M7 runbook. Adversarial review (`wf_87b22fa4`) folded a **blocker** (cloud`fetch` needs a dst ? reads silently returned `{}` on gcs/s3) + a poll-wait test gap + `PolicyTimeout`?clean rc 4`. Suite green, ruff+mypy clean. **M6** ? ? the real build/push/L4 run + `DEPLOY_REPORT` is M7 (owner, D17 budget).

169 - 2026-06-21 ? **M5 code landed** ([#286](https://github.com/Khelsutra/badminton-highlight-indexer/pull/286)): `backend/config/startup.py` ? per-profile fail/warn-fast startup checks (cloud-serving + non-cloud storage = the one fatal coherence error; GPU/creds mismatches = warnings, the runner healthcheck stays the authority), wired into the server (`_enforce_startup_or_exit` ? exit 1`) + the worker (`cmd_infer`/`cmd_train` ? rc 2`), default-OFF (probe gated behind an active profile ? the worker stays torch-free at `profile=""`). **L4 profile-parity test** (`tests/test_profile_parity.py`): same stub trajectory ? byte-identical windows + segment ranges across truly-local / cpu-mock / no-profile + worker-vs-in-process. Adversarial 4-lens review folded (`wf_34671ac1`). Suite green, ruff+mypy clean. **M5** ? ? the real-GPU truly-local runlog (L5) is owner-side (template pre-staged in `runlogs/`).

170 - 2026-06-21 ? **M9 hardening: exposure-safety boot posture** (alpha-shareable "tunnel-to-a-box" enabler): extended `backend/config/startup.py` with a **server-only exposure posture** (new `include_exposure=True` from `_enforce_startup_or_exit`; the worker stays unaffected ? it reads `--video-uri`, not untrusted tenant paths). Gated on `security.edge_auth_enabled` (**default-OFF** ? strict no-op, byte-identical to today): when the owner flips edge auth ON to expose the engine behind the CF-Access gate, three checks fire at **boot** ? `EXPOSED_NO_VIDEO_ROOT` (**fatal**: `allowed_video_root` jail unset ? a tenant could read an arbitrary local file`), `EXPOSED_EDGE_AUTH_INCOMPLETE` (**fatal**: CF team-domain/AUD missing ? lifts the current first-request 401 to boot), `EXPOSED_AUTH_EXTRA_MISSING` (**warn**: `PyJWT[crypto]` not importable ? would 500 per request). So a half-configured exposure fails fast instead of leaking/500-ing. +15 unit tests (default-OFF no-op, worker-ignores, both-fatals-compose-with-profile-checks, server-exit, WARN-only-no-exit). Adversarial 3-lens review (`wf_35b735a9`) confirmed default-OFF byte-identical + no bypass; folded a **whitespace-drift** (aligned `auth.resolve_tenant`` to `.strip()` team/AUD so "configured" means the same at boot + at request time) + named the config fields in the edge-auth-incomplete message. Pairs with the appliance **T3** network posture (raw `:8000` refused ? owner-run`). Suite green, ruff+mypy clean. **M9** ? the actual exposure flip + CF team/AUD/ingress + key rotation stay owner-side. (Deferred to the PR-2 runbook: document that `upload_dir`` must sit under `allowed_video_root``, and that you **MUST** set `edge_auth_enabled=True` when exposing ? these checks only fire then.)

171 - 2026-06-22 ? **M11 skeleton landed** (data-flywheel harness, D12): `backend/flywheel.py` (`capture_job` ? per-video `VideoWorkspace` via `backend.storage.put`;` `shadow_eval` ? experiment.compare`/`compare_unlabeled`;` `propose_promotion` ? promotion.promote_if_served_safe` with Gemini=control/owned=proposal;` `run_for_job` orchestrator`) + `FlywheelConfig` (**default-OFF**) + a gated `_maybe_run_flywheel` job-completion hook mirroring M8's cost hook`. **TWO** independent guards keep it INERT in production: `flywheel.enabled=False` (default) AND `consent_attested`` is a **FAKED hard-deny** (returns False for every job) ? so even enabled it captures **nothing** until the real consent schema lands. This is the PLUMBING only; it reuses the real workspace/eval/promotion/storage code so it works the day consent + the owned model + golden labels are wired. **NO DB migration** (the skeleton persists no consent). +13 tests (default-OFF no-op, enabled-but-consent-denied no-op, consented-capture writes the manifest, delegation of shadow-eval/promote, the hook is best-effort). Suite green, ruff+mypy clean. **M11** ? ? LIVE activation is owner/counsel (the consent ToS/DPA + schema ? the consent columns) + owner/GPU (the owned model + golden corpus for shadow-eval/promotion). ? the future M9-consent migration is now `v9`` (`v7` = `locale``, `v8`` = the compute-picker `compute`` column ? the schema ladder drifts as siblings land; check `PRAGMA user_version`` before authoring).

172 - 2026-06-22 ? **M12 backend landed** (`gdrive-api://` OAuth/API StorageBackend`, D14 Northstar): `backend/storage/gdrive_api.py` `GDriveApiStorage(StorageBackend)`` ? the headless **Drive-API** round-trip (read source + write reel back), distinct from the mount-only `gdrive://`` (left untouched). Path-keyed `gdrive-api://<path>` resolves each component ? a Drive file id via `files().list``; implements `fetch_to_local`` (`get_media().execute()`` bytes), `put`` (folder-creating `files().create`` + a `_media_body`` seam), `exists`/`stat`/`list`/`delete``. Wired the scheme into `ref.py`` (`_KNOWN`` + the path-keyed branch), `__init__._backend_class``, and the `StorageConfig.backend`/`input_backend`` Literals + a `gdrive_api`` settings dict (read via a **hyphen?underscore** lookup in `get_storage``). **All Google libs lazy-imported** (`import backend.storage`` never pulls them); a `service`` is **injectable** ? fully-offline tests via an in-memory fake Drive (round-trip + scheme + registry + the underscore-settings lookup);

`\_build\_service` (ADC over the Drive scope) raises a clear `storage-gdrive-api`-extra error when
libs/creds are absent. **NO DB migration.** Adversarial 2-lens review (`wf\_ae8c9b0f`). Suite
green (1466), ruff+mypy clean. **? M12 ? ? the LIVE per-user interactive OAuth app (consumer
"connect your Drive": client id/secret + refresh tokens) is owner-gated; the backend ships the
service-account/ADC build + the fake. Reads no secrets from config (creds via ADC only).**
173

4. user (2026-07-08T05:45:10.748Z)

```
total 116
drwxr-xr-x 1 avidu 197609    0 Jul  8 10:49 .
drwxr-xr-x 1 avidu 197609    0 Jun 25 16:24 ..
-rw-r--r-- 1 avidu 197609 4860 Jul  4 06:58 DEPLOY_REPORT_cloud-serving_2026-07-04.md
-rw-r--r-- 1 avidu 197609 9398 Jul  4 13:58 DEPLOY_REPORT_cloud-serving_a1-cuj_2026-07-04.md
-rw-r--r-- 1 avidu 197609 8687 Jul  4 17:11 DEPLOY_REPORT_cloud-serving_a1-e2e_2026-07-04.md
-rw-r--r-- 1 avidu 197609 8207 Jul  7 09:03 DEPLOY_REPORT_cloud-serving_alpha-
alliv1_2026-07-06.md
-rw-r--r-- 1 avidu 197609 8094 Jul  5 20:46 DEPLOY_REPORT_cloud-serving_cp-
rebuild_2026-07-05.md
-rw-r--r-- 1 avidu 197609 8474 Jul  7 19:17 DEPLOY_REPORT_cloud-serving_gpu-resident-
nvdec_2026-07-07.md
-rw-r--r-- 1 avidu 197609 21259 Jul  8 10:49 DEPLOY_REPORT_cloud-serving_validation-
latency_2026-07-07.md
-rw-r--r-- 1 avidu 197609 3609 Jul  7 23:11 README.md
-rw-r--r-- 1 avidu 197609 11336 Jul  7 23:11 RUNLOG_cloud-serving_exec-status-
fastfail_2026-07-07.md
-rw-r--r-- 1 avidu 197609 2698 Jun 25 16:24 RUNLOG_truly-local_TEMPLATE.md
```

5. user (2026-07-08T05:45:19.169Z)

```
1      # runlogs/ ? sample runs & deploy reports for posterity
2
3      This directory is the **durable record of real runs** of the compute-decoupled serving
profiles ?
4      the same convention proven by the archived
`DECOUPLED_COMPUTE/DEPLOY_REPORT_GCP_2026-06-20.md`. The
5      goal: **the next deploy is one-shot**, and anyone can trace exactly what was run, what
it cost, and
6      what came out.
7
8      ## When to add a file here
9      - Any **real-hardware run** (owner-side L5 in
[`../TESTING_STRATEGY.md`](../TESTING_STRATEGY.md)) of a profile.
10     - Every **default flip** that touches the core ships its Launch Report record here too
(per [[launch-report-convention]]).
11
12     ## Naming
13     `DEPLOY_REPORT_<profile>_<YYYY-MM-DD>.md` ? e.g. `DEPLOY_REPORT_truly-
local_2026-07-01.md`,
14     `DEPLOY_REPORT_cloud-serving_2026-07-01.md`.
15
16     ## Required contents (mirror the DECOUPLED deploy report)
17     1. **TL;DR** ? what ran, profile, ??? per stage (detect / window / stitch / serve).
18     2. **Resources** ? exact machine/VM/Cloud-Run config, region, image; for cloud, the
bucket + SA + budget.
19     3. **Exact commands** ? copy-pasteable (`worker infer`, the dispatch call, the
config/profile used).
20     4. **Run results** ? rally counts, **profile-parity numbers** (windows/segments vs the
other profile),
21     and the metering trio **`gpu_active_seconds` / `wall_seconds` / `bytes_out`** (even
before the M8 ledger lands).
22     5. **Reel smoke-test** ? `ffmpeg -f null -` full decode (exit 0) + a 5-frame contact
sheet visually
23     confirming real rallies (court/net/players mid-rally).
24     6. **Gotchas** ? conda-ToS, WASB-C3, GPU stockouts, font fallback, cold-start, NVDEC,
```

```

etc.
25      7. **Cost** ? measured/estimated; reconcile probe metrics vs the platform's per-
invocation billing.
26      8. **Teardown / final state** ? for cloud, the `$0` audit (`gcloud compute instances
list` + `disks list` empty); **DELETE, never just stop** ([[gcp-vm-always-delete]]).
27
28      ## DoD coupling
29      `COMPUTE_DECOUPLED_SERVING` is not DONE until **`truly-local` (M5)** and **`cloud-
serving` (M7)** each
30      have **?1 committed report here** (README §9).
31
32      ## Index (append as runs land)
33      | Date | Profile | Result | File |
34      |---|---|---|---|
35      | 2026-07-06 | cloud-serving | alpha-tester "All 1 v 1" timed out ? downsample fix
proven on L4 (720p30: 956 s, 60 rallies) | [DEPLOY_REPORT_cloud-serving_alpha-
all1v1_2026-07-06.md](DEPLOY_REPORT_cloud-serving_alpha-all1v1_2026-07-06.md) |
36      | 2026-07-07 | cloud-serving | GPU-resident NVDEC decode (#506/#507) shipped: image bump
+ cutover, cpu/nvdec smokes, GX010128 trajectory reproduced VM (26479?26481 vis), enabled +
canaried | [DEPLOY_REPORT_cloud-serving_gpu-resident-nvdec_2026-07-07.md](DEPLOY_REPORT_cloud-
serving_gpu-resident-nvdec_2026-07-07.md) |
37      | 2026-07-07 | cloud-serving | Crash-before-marker fast-fail (#515) LIVE smoke: real
`get_execution` classifier validated vs history (Succeeded+Failed); crash?**rc 5 in 47 s** (not
~65 min); healthy dispatch **rc 0 / 19 rallies**, probe returned None on all 9 running-polls |
[RUNLOG_cloud-serving_exec-status-fastfail_2026-07-07.md](RUNLOG_cloud-serving_exec-status-
fastfail_2026-07-07.md) |
38      | 2026-07-07 | cloud-serving | validation-latency closure (#512 shared-decode + #513
worker skip-reval): worker A/B on live image proves Video 1 FAIL?PASS (13 rallies) + re-val 169
s?70 s (2.43x) on L4; off-box #512 1.95?3.23x across codecs, zero decision drift; CP baseline
6m52s (live rev-00018 number owner-gated) | [DEPLOY_REPORT_cloud-serving_validation-
latency_2026-07-07.md](DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md) |
39

```

6. user (2026-07-08T05:45:19.573Z)

```

1      <!--
2      PROJECT DOC TEMPLATE ? copy this file to docs/<YOUR_PLAN>.md and fill it in.
3
4      A "tracked project doc" is how this repo carries substantial design/project work: it can
be
5      ITERATED on, its progress TRACKED, and it is MARKED DONE against explicit deliverables,
then
6      archived. Standardizes the emergent pattern in docs/RALLY_QUALITY_RESEARCH.md §7 and
docs/GOLDEN_REGRESSION_FIXTURES.md.
7
8
9      Rules of thumb:
10     - §7 (the progress tracker) is the SOURCE OF TRUTH ? one row per independently-shippable
PR.
11     - Keep it HONEST: tag each claim [verified] (checked against code), [researched] (needs
sign-off),
12       or [design] (proposed). The doc reflects real shipped state, never aspirational.
13     - It POINTS to detail (code anchors, research artifacts), it does not duplicate it.
14     - Sections marked OPTIONAL are included only when relevant (e.g. Foundation/Threat-model
for
15       research- or security/privacy-heavy work). Delete sections that do not apply.
16     - Move to docs/archives/ once Status = DONE (per the archive policy: only terminal-state
work).
17     -->
18
19     # <Project / Plan Title>
20
21     > **Status:** `DRAFT` ? `<owner-gated? | in progress>` . **Owner:** `<name>` .
**Created:** `<YYYY-MM-DD>` . **Last updated:** `<YYYY-MM-DD>`
22     > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE)

```

```

23      > **Tracking anchors:** §7 progress tracker is the source of truth; indexed in
`docs/README.md`; pointer in memory `current-state` (and `docs/NEXT_STEPS.md` once work starts).
24      > **Relation to existing docs:** `<extends | supersedes | peer-of ?>`
25      > **Honesty note:** mark each claim `[verified]` / `[researched]` / `[design]`. Not
legal/financial advice where applicable.
26
27      ---
28
29      ## 0. TL;DR
30      <2?5 sentences: the goal, the approach, and the single most important caveat. Call out
any finding that is already actionable on its own.>
31
32      ## 1. Problem & goal
33      <Why now; the concrete user-facing outcome; what "good" looks like. Link the code/docs a
reader must read to get current.>
34
35      ## 2. Decisions locked
36      <Table ? capture answered questions so they are not relitigated.>
37
38      | # | Decision | Source / date | Implication |
39      |---|-----|-----|-----|
40      | D1 | ? | ? | ? |
41
42      ## 3. Foundation ? research / prior art *(OPTIONAL)*
43      <Include when the plan rests on external facts: cited research, legal/compliance,
benchmarks, prior-art survey. Mark researched-vs-verified; list any sign-off gates this
creates.>
44
45      ## 4. Design / architecture
46      <The actual plan. Data contracts, diagrams, and a REUSE MAP ? name the existing
modules/code this builds on so it is clear what is new vs reused.>
47
48      ## 5. Threat model / risk table *(OPTIONAL ? security / privacy / compliance work)*
49      <Who/what is protected vs NOT; residual risks stated plainly.>
50
51      ## 6. Honest limits ? what this does NOT do
52      <The false-confidence guardrail: what a green suite / a DONE status does NOT prove.>
53
54      ## 7. Deliverables & progress tracker ? **source of truth**
55      Legend: ? Todo · ? In progress · ? Done · ? Blocked/gated. **One small PR per row**,
branched from `origin/master`, no stacking. Update the row (status + PR link) as work lands.
56
57      | ID | Deliverable | Depends on | Gated? | Status | PR |
58      |---|-----|-----|-----|-----|----|
59      | P0 | ? | ? | No | ? | ? |
60
61      ## 8. Open questions ? owner / external *(split blocking-gates from nice-to-knows; name
who decides)*
62      1. ?
63
64      ## 9. Definition of done
65      <Explicit acceptance checklist. When all ticked ? flip `Status ? DONE` and archive.>
66      - [ ] ?
67
68      ## 10. References
69      **Internal code anchors `[verified]`: **`<paths>`. **Internal docs:** **`<docs>`.
**External:** **`<cited sources>`.
70
71      ---
72
73      ### Changelog
74      - **`<YYYY-MM-DD>` ? **`<what changed>`.
75

```

```

1      # DEPLOY REPORT ? cloud-serving, validation-latency closure (#512 shared-decode + #513
worker skip-revalidation), 2026-07-07
2
3      > Durable **closure record** for the two validation-latency fixes now live on the
`cloud-serving` path: ##512** (`perf(validators)`): decode `scene_cut`/`blur`/`relevance` in ONE
shared pass instead of ~125 random per-sample seeks across 3 captures) and ##513**
(`fix(serving)`): the worker skips its redundant re-validation when the control-plane already
validated + dispatched, via `--skip-validation`). Both shipped in the live image `rally-
worker@sha256:8045f54e?` (built from `origin/master` @ `f5c9dd9`; #512 `20862d7`, #513 `23c1868`
are ancestors). Measures the before/after on 3 clips spread across codecs/resolutions, all
against the real serving infra. Project `gen-lang-client-0306026826`, region `asia-southeast1`.
Honest-limits (incl. the owner-gated live CUJ) in $6.
4
5      ## 1. TL;DR
6
7      | Stage | Result |
8      |---|---|
9      | Live image carries #512+#513 | ? control-plane rev `00018-qvz` + worker Job both on
`sha256:8045f54e?` (= `f5c9dd9`, ancestors `20862d7`/`23c1868` verified) |
10     | Baseline captured (pre-fix) | ? from prod logs on rev `00017-qbd` (image `09d2edf7` =
`66756ad`, pre-#512/#513) ? no re-run needed |
11     | ##513 worker skip-revalidation** | ? real dispatch-style exec: Video 1 validation
**SKIPPED (0 s)**, runs to **SUCCESS / 13 rallies** ? the blur-reject that failed the live CUJ
is **gone** |
12     | ##512 shared-decode speedup (real L4)** | ? Video 1 re-validation **169.3 s ? 69.7 s
(2.43x)** on identical worker hardware, decision byte-identical (sharpness 11.9 < 20.0 both) |
13     | ##512 speedup (off-box A/B, breadth)** | ? 1080p H.264 **1.95x**, 4K HEVC **3.10x**,
5.3K 10-bit HEVC **3.23x**; every validator decision unchanged OLD?NEW |
14     | Live control-plane "after" number | ? **live-corroborated 2026-07-08** ($9): rev
`00019` validated a fresh 1.83 GiB / 174 s clip in **183 s** (in the 2?3 min band) + reel
compiled end-to-end. Video-1's exact 412 s?~170 s A/B stays a projection (the live CUJ used a
different, heavier clip). |
15
16     **Bottom line:** the two fixes do exactly what they claim on live serving hardware.
##513** turns the soft 5.3K clip from a **hard FAIL** (worker re-validated at its default
`min_blur=20`, rejected sharpness 11.9, `rc2`, no reel) into a **PASS with 13 rallies** by
trusting the control-plane's gate. ##512** cuts the frame-sampling validators ~**2?3x** (2.43x
measured on the L4; 1.95?3.23x off-box across codecs) with **zero decision drift** OLD?NEW. The
only number still owner-pending is the live control-plane validation wall on rev 00018 ? its
components are measured and it projects to ~2?3 min from the ~6m52s baseline.
17
18     ## 2. Resources
19     - **Control-plane:** `rally-control-plane` Cloud Run **service**, rev **`rally-control-
plane-00018-qvz`** (100% traffic), image `asia-southeast1-docker.pkg.dev/gen-lang-
client-0306026826/rally/rally-worker@sha256:8045f54ea9fcbf?` (built from `f5c9dd9`),
`DEPLOYMENT_PROFILE=cloud-serving`, CPU-only **2 vCPU**, edge-auth ON (`edge_auth_enabled=true`,
`cf_team_domain=khelsutra.cloudflareaccess.com`).
20     - **Worker:** `rally-worker` Cloud Run **Job**, same image `sha256:8045f54e?`, 8 vCPU /
32 GiB / 1x NVIDIA **L4**, `WASB_WEIGHTS=/gcs/models/wasb_badminton_best.pth.tar`,
`WASB_DECODE_BACKEND=nvdec_resident`, GCSFuse mount of `gs://khelsutra-rally-serving`.
21     - **Baseline (pre-fix) surface:** control-plane rev **`00017-qbd`** / worker image
`sha256:09d2edf7?` (= `66756ad` = #507+#508+#509, both **pre-#512/#513**) ? the source of the
prod baseline logs below (2026-07-07, before the 16:15 rev-00018 cutover).
22     - **Corpus:** Video 1 `gs://khelsutra-rally-
serving/videos/uploads/702636eed76145e88aa59c3000c7e415/Video 1 - Badminton.MP4` (1.41 GiB,
5312x2988 yuv420p10le HEVC, 3019 frames, MD5 `8fe6b719603dd7a7b2b6ef543139a4a3`);
`gs://khelsutra-rally-corpus/videos/golden_1080p/Boxhill_Doubles.mp4` (1080p H.264);
`gs://khelsutra-rally-corpus/videos/golden/GX010128.mp4` (2.18 GiB, 4K HEVC, 45298 frames).
23
24     ## 3. Exact commands
25     ```bash
26     # --- Provenance: prove the live image carries both fixes ---
27     gcloud run services describe rally-control-plane --region asia-southeast1 \
28     --format="value(status.latestReadyRevisionName,
spec.template.spec.containers[0].image)" # 00018-qvz ?@sha256:8045f54e?

```

```

29     git merge-base --is-ancestor 20862d7 f5c9dd9 && echo "#512 in live image"    # yes
30     git merge-base --is-ancestor 23c1868 f5c9dd9 && echo "#513 in live image"    # yes
31
32     # --- #513 proof: worker A/B on the LIVE image (Video 1 already staged in the serving
33     # bucket) ---
34     V1="gs://khelsutra-rally-serving/videos/uploads/702636eed76145e88aa59c3000c7e415/Video 1
35     - Badminton.MP4"
36     # Run A ? WITH --skip-validation (exactly what orchestration._maybe_dispatch_remote
37     # emits on dispatch):
38     gcloud run jobs execute rally-worker --region asia-southeast1 --async \
39     --args="^@@^infer@@--video-uri=$V1@@--out-uri=gs://khelsutra-rally-serving/_valclose/_
40     valclose/v1-skipval@@--segmenter=wasb_hybrid@@--skip-validation@@--
41     set=deployment.compute_target=local_gpu@@--set=indexing.detector_impl=native@@--
42     set=indexing.skip_ai_handoff=true@@--set=indexing.wasb.batch_size=64@@--
43     set=indexing.wasb.prefetch_batches=4@@--set=indexing.wasb.timeout_sec=0"
44     # Run B ? WITHOUT --skip-validation (worker re-validates with the NEW #512 validators ?
45     # still blur-rejects):
46     gcloud run jobs execute rally-worker --region asia-southeast1 --async \
47     --args="^@@^infer@@--video-uri=$V1@@--out-uri=gs://khelsutra-rally-serving/_valclose/v
48     1-reval@@--segmenter=wasb_hybrid@@--set=deployment.compute_target=local_gpu@@--
49     set=indexing.detector_impl=native@@--set=indexing.skip_ai_handoff=true"
50
51     # --- #512 timing: bracket "pulled ? -> in.mp4" ? "validation FAILED/SKIPPED" ?
52     "Phase 2 (WASB) completed" ---
53     gcloud logging read 'resource.type="cloud_run_job" AND
54     labels."run.googleapis.com/execution_name"="rally-worker-<exec>" \
55     --freshness 1d --format="value(timestamp, textPayload)"
56
57     # --- #512 off-box A/B (isolated clones; run_all decodes the SAME sample indices in both
58     # trees) ---
59     git clone --local --no-hardlinks <repo> repo_old && git -C repo_old checkout 23c1868 #
60     pre-#512 (per-sample seeks)
61     git -C repo_old worktree add --detach repo_new f5c9dd9
62     # live image src (shared pass)
63     python valbench.py <repo_old|repo_new> <video> <label> 2 # times
64     registry.run_all(video, load_config())
65     ```
66
67     ## 4. Run results
68
69     ### 4a. Video 1 (5312x2988 10-bit HEVC, 3019 frames) ? the clip that motivated #512/#513
70
71     | Metric | BEFORE ? pre-fix (rev `00017`, img `09d2edf7`) | AFTER ? #512+#513 live (rev
72     `00018`, img `8045f54e`) |
73     |---|---|---|
74     | Control-plane validation | **411?412 s** (6m52s) ? passed (baked `min_blur=8.0`) |
75     **projected ~170 s (2?3 min)**; **live-corroborated 2026-07-08 ? 183 s** on rev `00019` ($9 ? a
76     fresh 1.83 GiB clip, not a Video-1 re-run) |
77     | Worker (re-)validation | **169.3 s** ? ? **BLUR-REJECT** (sharpness 11.9 < default
78     20.0) | **0 s ? SKIPPED** (`--skip-validation`, #513) |
79     | &nbsp;&nbsp;&nbsp;? counterfactual: worker re-val if NOT skipped | *(169.3 s, OLD
80     validators)* | **69.7 s** (NEW #512 validators, **2.43***, same 11.9<20 decision) |
81     | End-to-end | ~10m31s wall ? **FAILED** (0 rallies, no reel) | worker ? **SUCCESS** in
82     ~2m38s exec ? **13 rallies**; **reel live-confirmed 2026-07-08** ($9 ? fresh-clip CUJ compiled
83     an 11-clip reel end-to-end) |
84     | Rallies produced | **0** (rejected before detection) | **13** (`report.json`
85     `status=success`, 13 segments) |
86     | Pass / fail | ? **FAIL** | ? **PASS** |
87
88     **Prod baseline evidence (rev 00017, 2026-07-07, unmodified):** control-plane `13:49:18
89     Running validations ? 13:56:10 Video passed checks` = 411.4 s (and `14:03:56 ? 14:10:48` = 411.9
90     s). Worker exec `rally-worker-2lbz2` (dispatched by the 14:03 run, OLD image): `pulled 14:11:35
91     ? validation FAILED: blur_check ? sharpness 11.9 below 20.0 @14:14:24 ? exit(2)` ? a
92     **successful GPU dispatch that failed after the fact**, the exact CUJ #513 fixes.
93
94     65

```

```

66      **After evidence (live image `8045f54e`, 2026-07-07):**
67      - `rally-worker-k29rk` (**Run A, --skip-validation**): `pulled 17:09:24 ? [worker]
validation SKIPPED (--skip-validation) 17:09:27 ? native WASB (torch 2.6.0+cu124, cuda True) ?
Phase 2 (WASB) completed in 100.9 s, 13 candidate windows (1510/3019 visible) ? emitted 13
rallies ? exit(0)`. Exec wall 157.7 s; `gpu_active`?100.9 s (WASB);
`bytes_out`=`report.json`+`intervals.csv`. Output
`gs://?/_valclose/v1-skipval/outputs/8fe6b719?/`, `status=success`, 13 segments.
68      - `rally-worker-t8llx` (**Run B, re-validation on the NEW image**): `pulled 17:09:29 ?
validation FAILED: blur_check ? sharpness 11.9 below 20.0 @17:10:39 ? exit(2)`. Re-val wall
**69.7 s** vs the OLD image's 169.3 s = **2.43x on identical L4 hardware**, and the **decision
is byte-identical** (11.9 < 20.0) ? #512 is pure speed, zero behavior change. (This run also
shows *why* #513 is needed: absent the skip, the worker's default `min_blur=20` still rejects
the clip the control-plane passed at 8.0.)
69
70      ### 4b. Off-box #512 A/B ? codec/resolution breadth (16-core local box, cv2 4.13.0,
warm, `run_all`)
71
72      Isolated clones at `23c1868` (OLD, per-sample seeks) vs `f5c9dd9` (NEW, shared pass);
identical sample indices ? a fair A/B. **Every validator decision is identical OLD?NEW** on
every clip (no PASS/FAIL drift from the fix).
73
74      | Clip | codec / res | OLD (warm) | NEW (warm) | Speedup | decisions OLD?NEW |
75      |---|---|---|---|---|---|
76      | Boxhill_Doubles | 1080p H.264 | 10.91 s | 5.60 s | **1.95x** | 6/6 identical (all
PASS) |
77      | GX010128 | 4K HEVC (45298 fr) | 35.76 s | 11.54 s | **3.10x** | 6/6 identical (all
PASS) |
78      | Video 1 | 5.3K 10-bit HEVC | 173.20 s | 53.62 s | **3.23x** | 6/6 identical (all
PASS?) |
79
80      ? On this Windows/cv2 box the 10-bit-HEVC decode yields a Laplacian variance **above**
20 so blur PASSES off-box, whereas the Linux serving stack measures **11.9** and rejects. That
platform decode-stack difference is **orthogonal to #512** (which changes only *which frames are
sampled*, identically in both trees); the authoritative blur verdict and the real speedup come
from the L4/prod ($4a). Off-box contributes only the OLD?NEW **ratio**, which is decode-stack-
independent.
81
82      The win grows with per-frame decode cost (random-seek re-decode is what #512 removes),
and it **triangulates** with the real-L4 Video 1 A/B (2.43x): the off-box ratios bound the
control-plane projection.
83
84      ### 4c. Control-plane projection (2-vCPU)
85      Baseline **~412 s**. Applying the L4-measured **2.43x** ? **~170 s (~2m50s)**; the off-
box ratios (1.95?3.23x) bracket **~2?3 min**. Matches the pre-fix design estimate (~6.5 min ? ~2
min). Exact live number is owner-gated ($6).
86
87      ## 5. Rollout state
88      - **Already deployed ? this is a measurement/closure record, not a deploy.** #512+#513
went live with the `8045f54e` image cutover (control-plane rev `00018-qvz`, worker Job same
digest) on 2026-07-07. No infra change was made for this report.
89      - **Rollback (pre-existing levers):** control-plane `gcloud run services update-traffic
rally-control-plane --region asia-southeast1 --to-revisions rally-control-plane-00017-qbd=100`;
worker `gcloud run jobs update rally-worker --region asia-southeast1 --image ?/rally-
worker@sha256:09d2edf7?`. Both revert to the pre-#512/#513 image.
90      - **Scratch outputs** written under `gs://khehsutra-rally-serving/_valclose/` (isolated
prefix; safe to delete).
91
92      ## 6. Honest limits / not-verified
93      - **Live end-to-end CUJ is owner-gated.** ? **CLOSED 2026-07-08 ? see $9** (owner drove
one live SPA upload on rev `00019`; both ? cells now measured). Original limit, retained for
context: `POST /api/process` on rev 00018 requires a Cloudflare-Access JWT
(`edge_auth_enabled=true`; a direct call returns **401** `missing Cf-Access-Jwt-Assertion
header`). So the two ? cells ? the **live control-plane validation wall on rev 00018** and the
**stitched reel** ? need the owner to drive one upload (or hand over a CF-Access JWT / service
token). Everything else here is measured on the real serving infra. **Recipe to close them:**

```

upload "Video 1 - Badminton.MP4" via the SPA (compute=cloud), then `gcloud logging read '?service\_name="rally-control-plane"?("Running validations" OR "passed checks")'` for the wall, and confirm the dispatched worker exec logs `validation SKIPPED` + emits rallies + the reel compiles. Expect CP validation ~2?3 min and a PASS.

94 - ****#513's worker win is shown via a dispatch-faithful direct exec**, not a control-plane-initiated dispatch (same auth gate). The `--skip-validation` flag, its emission in `orchestration.\_maybe\_dispatch\_remote`, and its handling in `worker.cmd\_infer` are unit-tested on master; Run A exercises the exact worker code path with the exact flag. ? ****UPDATE 2026-07-08 (\$9): now also proven on a genuine control-plane-initiated dispatch**** (`rally-worker-64bhj`): `\_maybe\_dispatch\_remote` emitted `--skip-validation` on a live SPA run and the worker honored it (`validation SKIPPED`, 0 re-validation).**

95 - ****Rallies, not a reel.**** Run A used `skip\_ai\_handoff=true` (secret-free) so the 13 "rallies" are local candidate windows (no Gemini refine, no stitched mp4) ? this validates decode+detect+windowing end-to-end and that the clip ****passes****; the reel stitch is the owner-CUJ step above. The 13 segments are its inputs.

96 - ****Off-box absolute times are from a Windows 16-core box (cv2/FFmpeg), not the 2-vCPU Linux control-plane**** ? so the ****ratios**** transfer (the change is algorithmic: seeks?shared forward-hybrid pass), the absolutes don't, and (per ?) the 10-bit blur ***verdict*** can differ; the control-plane absolute is anchored to the prod 412 s baseline, not the local numbers. Video 1 off-box is a cross-check of the authoritative real-L4 2.43x.

97 - ****#515**** (crash-before-marker fast-fail) is NOT in this image (`ef41473` post-dates `f5c9dd9`); it's unrelated to validation latency and out of scope for this closure.

98

99 **## 7. Cost**

100 - 2 L4 worker execs (Run A 157.7 s + Run B 109.6 s ? 268 s of L4 time) ? ****~\$0.5?0.7****. Off-box A/B: ****\$0**** (local). No VM, no image build. GCS read for the 3 downloaded clips (~5.1 GiB) is intra-project.

101

102 **## 8. Closure verdict**

103 ? ****CLOSED ? #512 and #513 are proven on live cloud-serving hardware.**** #513 converts the motivating soft-focus 5.3K clip from a post-dispatch FAIL into a 13-rally PASS by eliminating the worker's redundant re-validation (169 s + wrongful blur-reject ? 0 s); #512 speeds the frame-sampling validators ****2.43x on the real L4**** and ****1.95?3.23x off-box across codecs****, with ****zero decision drift**** OLD?NEW on every clip. The last owner-gated residual is now ****closed live 2026-07-08 (\$9)****: an owner-driven SPA CUJ on rev `00019` validated a fresh 1.83 GiB / 174 s clip in ****183 s**** (in the projected 2?3 min band) and compiled an 11-clip reel end-to-end ? with the worker `validation SKIPPED` on a genuine control-plane dispatch. (The Video-1-specific 412 s?~170 s A/B stays a projection by design ? the live CUJ used a different, heavier clip.) Validation-latency work is done.

104

105 **## 9. Update (2026-07-08) ? live SPA CUJ closes the two owner-gated cells (on rev `00019`)**

106

107 The two ? cells in \$4a (the ****live control-plane validation wall**** and the ****end-to-end reel****) were owner-gated on the Cloudflare-Access `POST /api/process`. On ****2026-07-08**** the owner drove one real SPA upload (compute=****Cloud****) end-to-end, so both are now measured on live serving infra. Mirrors the gpu-resident report's "\$5 Update" append ? the body above is unchanged; this section records what the live run added.

108

109 ****Two honest deltas from the recipe**** (state moved after this report's 2026-07-07 image cutover):

110 1. ****Live rev rolled `00018` ? `00019`.**** Traffic is 100% on ****`rally-control-plane-00019-96f`****, image `rally-worker@sha256:c6393d27b4ae?` (tag `:latest`), deployed 2026-07-07T17:06Z ? ~51 min after the `00018`/`8045f54e` cutover \$2 documents; the worker Job is on the same `:latest`=`c6393d27` digest (both surfaces match). It's an undocumented forward-bump, but #512 (`20862d7`) and #513 (`23c1868`) are ancestors of `origin/master`, `f5c9dd9` (the `00018` src) is an ancestor of master, and the run ****empirically**** shows both fixes live (one-pass validation wall + worker `validation SKIPPED`). Measured on what is actually serving.

111 2. ****A fresh clip, not Video 1.**** The upload was ****"First CUJ.MP4"***** ? ****1.83 GiB**** (1,960,043,219 B), source duration ****174.17 s**** (?84 Mbps, 4K-class ? ***heavier*** than Video 1's 1.41 GiB), video_id `2f31b45e?`. So the live wall ****corroborates**** the Video-1 projection on a comparably-heavy clip; it is ****not**** a re-measurement of Video 1's 412 s?~170 s A/B (a different clip has a different absolute wall).

112

113 ****Full phase breakdown (UTC 2026-07-08, rev `00019`):****


```

114
115 | # | Phase | Window | Wall | Machine | Notes |
116 |---|---|---|---|---|---|
117 | 1 | Upload (client?server, 80-part multipart) | 04:25:56 ? 04:28:51 | **~2m55s** |
client uplink | 1.83 GiB @ ~84 Mbps ? owner bandwidth, not infra |
118 | 2 | Accept ? validation start (stage) | 04:28:52 ? 04:29:22 | ~30 s | control-plane |
stage upload to serving bucket |
119 | 3 | **Validation (#512 shared-decode)** ? | 04:29:22 ? 04:32:25 | **183 s (3m03s)** |
control-plane 2 vCPU | one-pass scene_cut/blur/relevance; PASSED |
120 | 4 | Dispatch ? worker ready (cold start) | 04:32:25 ? 04:33:20 | ~55 s | Cloud Run Job
| image pull + GCSFuse mount + torch/cuda init |
121 | 5 | **Worker skip-validation (#513)** ? | 04:33:20 ? 04:33:24 | ~4 s | L4 |
`validation SKIPPED`, 0 re-validation |
122 | 6 | Detection / segment (WASB) | 04:33:24 ? 04:36:20 | **161.6 s** | L4 (cuda) | 15
candidate windows, 2561 visible pts |
123 | 7 | Ingest + finalize | 04:36:20 ? 04:36:24 | ~4 s | control-plane | intervals?DB,
`status=success` |
124 | ? | (user reviews rallies, clicks Compile) | 04:36:24 ? 04:37:26 | ~62 s idle | ? |
human-in-the-loop |
125 | 8 | **Compile / stitch** ? slow | 04:37:26 ? 04:50:21 | **~12m55s** | control-plane 2
vCPU | 11 clips, ~67 s/clip 4K re-encode+watermark; concat ~2 s; mirror ~11 s |
126
127 Server pipeline (process?success, excl. upload+idle): **7m32s**. Full CUJ (upload?reel):
**~24m25s** wall.
128
129 **Cell 1 ? control-plane validation AFTER (live):** **183 s (3m03s)** on rev `00019`,
PASSED ? the #512 shared-decode validators on the real 2-vCPU serving path, in the projected
**2?3 min** band on a clip *heavier* than Video 1 (pre-#512 per-sample seeks were the
~6m52s-class baseline). Evidence: `Running validations 04:29:22.789 ? Video passed checks
04:32:25.780`.
130
131 **Cell 2 ? end-to-end reel (live):** `/api/compile` stitched **11 of the 15** detected
rallies (0.5 s pad, per-clip branding watermark) from the 174.17 s source ? **`gs://khehsutra-
rally-serving/highlights/2f31b45e?/first-cuj-ks-highlight-reel.mp4`** (**581.84 MiB**, saved
04:50:10 "Dynamic highlight reel saved (581.8 MB)", mirrored 04:50:21, served via
`/api/highlights` 307). **Reel ? produced end-to-end on the live control-plane** ? source fetch
? per-clip trim+watermark ? concat ? mirror-upload, all exercised.
132
133 **Bonus ? #513 proof upgraded** from the $6 dispatch-faithful *direct exec* to a genuine
**control-plane-initiated dispatch** (`rally-worker-64bhj`): `_maybe_dispatch_remote` emitted
`--skip-validation` on a live SPA run and the worker honored it (0 re-validation).
134
135 *** Finding ? stitch is the scaling bottleneck (follow-ups filed).** The compile ran
**~13 min on the CPU-only 2-vCPU control-plane** (~45?90 s per ~7 s clip: the watermark forces a
full 4K re-encode, no NVENC). With validation (183 s) *also* on the control-plane and serialized
by `_LOCAL_COMPUTE_LANE`, one run pins the control-plane for ~16 min of heavy CPU ? it won't
scale past a few concurrent users. Filed: **$521** (move compile/stitch to the L4 worker + make
phase?machine placement config-driven) and **$520** (step-level debug tracing, server+UI, to a
low-latency sink ? this breakdown had to be hand-correlated across 3 log streams).
136
137 **Cost (heuristic):** one L4 detect exec (`rally-worker-64bhj`, ~222 s) ? **$0.15?0.55**
(rate-dependent, dominant) + control-plane CPU for validation+stitch (~960 s @ 2 vCPU) ?
**~$0.05** + reel egress (581 MB, if downloaded to internet) ? **~$0.07** + storage/ops **<
$0.02** + Gemini **$0** (skip_ai_handoff). *** $0.3?0.7 total**, ~one L4 exec as projected ? but
the stitch's real cost is **concurrency** (~16 min of control-plane heavy CPU), not dollars.
138
139 **Net:** both owner-gated cells are now live-measured on rev `00019` ? validation wall
**183 s** (corroborating the 2?3 min projection) and the reel compiles end-to-end. #512 one-pass
validation and #513 skip-validation-on-real-dispatch are now both directly observed in prod. The
Video-1 412 s?~170 s A/B remains a projection by design (different, heavier clip).
140

```

8. user (2026-07-08T05:45:27.116Z)

```

2
3     > **Status:** ? **LIVING** ? the single board for doc **lifecycle + archival due-
diligence**. **Last full review:** 2026-07-02.
4     > **Update cadence:** whenever a doc changes lifecycle state (created / shipped a
milestone / delivered / archived), and re-review at each clean-slate.
5     > **Why this exists:** so we can (a) **review completeness** at a glance, (b) **update
it as work finishes** to prove we're done, and (c) **never archive a foundational doc without
honestly updating it first** (owner directive, 2026-06-21).
6
7     ---
8
9     ## 0. How to use this
10    - **Reviewing now?** Read §4 (the open findings/actions) ? that's what needs attention.
§3 is the full register.
11    - **Marking something complete?** Update the doc's own **§7 progress tracker** (the
source of truth for *its* completeness), then flip its row here, then ? if terminal ? run the §2
archival checklist.
12    - **This register does NOT duplicate completeness detail** ? it indexes lifecycle +
archive-readiness and points to each tracked doc's §7.
13
14    ## 1. Lifecycle legend
15    `DRAFT ? IN PROGRESS ? DONE/DELIVERED ? ARCHIVED`. Plus: **LIVING** (permanent
reference, never "done"), **PROPOSAL/DESIGN** (owner-gated, not started), **REPORT** (a terminal
record of a run/launch), **CONVENTION** (a standing rule/guide).
16
17    ## 2. Archival due-diligence checklist ? *the repeatable process (do NOT just `git
mv`)*
18    Before a foundational/tracked doc moves to `docs/archives/`:
19    1. **VERIFY claims against the code** ? no stale assertions; tag load-bearing facts
`[verified]`. (A doc that says "nothing built" when code shipped is *not* archive-ready ? fix it
first.)
20    2. **Flip Status ? DELIVERED / DONE / REJECTED + a completion note** ? what shipped,
where (PRs, paths, artifacts), and what was **deferred / carried forward**.
21    3. **Update ALL inbound references** ? `docs/README.md` index, `CLAUDE.md`,
`docs/NEXT_STEPS.md`, and any cross-doc links ? the new path + status.
22    4. **Keep the lineage** ? ADR forward/back-links + "supersedes / superseded-by" so the
evolution stays traceable.
23    5. **`git mv` the whole dir/file ? `docs/archives/past_projects/`** (terminal-state
ONLY; live docs stay at `docs/` top level). Relative links inside become historical (note this
in the doc's header).
24    6. **Record the move** ? flip the row here to §3e + a line in memory `current-state`.
25
26    > **Worked example (the template):** **DECOUPLED_COMPUTE ? #270** ? verified M0?M2+M3.5
delivered on GCP, flipped ??? DELIVERED with a completion note (incl. carried-forward
WASB-C3/DPA/`md5`), reindexed `docs/README` + `past_projects/README`, kept the ADR-010?014
lineage, then `git mv` to `docs/archives/past_projects/DECOUPLED_COMPUTE/`.
27
28    ## 3. The register
29
30    ### 3a. Tracked projects ? lifecycle docs (archive when DONE; each doc's §7 is the
completeness source-of-truth)
31    | Doc | Status | Archive-ready? | Next action |
32    |---|---|---|---|
33    | `ALPHA_LAUNCH_READINESS.md` | **DRAFT** ? cross-repo alpha readiness tracker created
2026-07-03 after the golden corpus reached n=15; separates hosted product-alpha work from owned-
model/commercial rally-detection gates. A0?A9 DONE; A1 live (Cloud Run, e2e runlog); **A13
review DELIVERED 2026-07-04** (`ALPHA_LAUNCH_REVIEW_2026-07-04.md`) | No | Owner: record G0/N0-
G0 (§10) + close/defer A10/A11/A12; fix khelsutra #154 copy blockers; run the 5-min browser-CUJ
checklist (review §1.1) |
34    | `COMPUTE_DECOUPLED_SERVING/` | **IN PROGRESS** ? M0 approved + **M1?M9 LARGELY
SHIPPED** (#279?#290; M5 #286, M6 #287, M8 #288, M9-auth #290 all merged, default-OFF); Gen-0
weights pushed (`weights/0.1.0-l4/`) | No | M11 flywheel + M12 gdrive-api (owner residual:
counsel/prices/OAuth/real-run within the D17 cap) |
35    | `RALLY_QUALITY_RESEARCH.md` | Foundation, **tracked & ACTIVE** (§7) | No | Drives the
R-series; complete when §7.1 single quality number is attributed |

```

36 | `R2\_EVAL\_METRIC\_DESIGN.md` | IMPLEMENTED (augmenting) | No (gates R4 + dual-eval; ablation pending) | Promote-or-reject toLF1 as the gate (ablation) ? then archive with the R-cluster |

37 | `R1\_MILESTONE\_LAUNCH\_REPORT.md` | LAUNCHED (#204) ? terminal record | Soon (with the R-cluster) | Keep with R-series for narrative; archive when the track wraps |

38 | `RALLY\_DETECTION\_GAP\_CLOSURE.md` | DRAFT (owner-gated) | No | G1?G3 levers; part of the R/quality cluster |

39 | `RALLY\_DETECTION\_FIXES\_PLAN.md` | **DONE** ? findings A/B/C resolved; P2b promoted `inactivity\_temporal\_density` default-ON (#396) | With the R-cluster (deferred per its header; gate = R2 toLF1 ablation + \$7.1 number) | Archive with the R-series when the track wraps |

40 | `RALLY\_DETECTION\_CODE\_REVIEW.md` | Terminal snapshot ? findings A (#394), B (#394 ? promoted #396), C (`sdc#50`) all resolved | With the R-cluster | Archive with the R-series |

41 | `PER\_VIDEO\_WORKSPACE.md` | DRAFT, IN PROGRESS (P0 = #264) | No | ? supersede-reconcile (finding #1) DONE; P3 de-hardcode base ? via #282; remaining P1?P6 (P1 v6 migration coordinates with M8's v6 cost-ledger ? version bump authored once) |

42 | `MULTI\_SIGNAL\_FUSION\_PLAN.md` | P1 DONE; P2+ owner-gated | No | P3/P4 (learned fusion, audio) ? needs GPU golden features |

43 | `EXPERIMENT\_HARNESS.md` | P1 IMPLEMENTED | No | A/B + promotion gate; supports the quality track |

44 | `PROMINENT\_COURT\_DETECTION.md` | DRAFT (owner-gated) | No | C-series; multi-track selection design |

45 | `EXECUTION\_POLICY.md` | P1 SHIPPED; P3b next | No | Async job policy; reused by Cloud Run dispatch |

46 | `GOLDEN\_REGRESSION\_FIXTURES.md` | IN PROGRESS (owner-gated) | No | Golden eval fixtures |

47 | **Data layer** ? `docs/data\_pipeline/` (`GOLDEN\_SET\_IMPLEMENTATION` . `GOLDEN\_SET\_PHASE1\_VIDEOS` . `GOLDEN\_VIDEOS` . `VIDEO\_RECORDING\_GUIDELINES` . `GOLDEN\_DATA\_SHARING`) | LIVING (data, **no-ML**) | n/a | Relocated 2026-06-21 to the isolated data pillar; tracked via `data\_pipeline/README.md`. Roora *vendor* docs archived ? `archives/roora-vendor/`. |

48 | `archives/past\_projects/AUDIT\_REPORT\_khelsutra-guru\_2026-07-02-COMPLETED-2026-07-02.md` ***(umbrella)** | **? COMPLETE / ARCHIVED (2026-07-02)** ? the five-repo code audit + folded-in *Remediation tracker* (P0?P28 all closed) + absorbed prior audits. All non-owner-gated residuals landed; owner-gated/future items filed as issues (#443/#444/#445, + #301/#327/#332/#384). | Archived | Owner confirmed archival 2026-07-02; moved to `archives/past\_projects/` per §2. Residuals now live as GitHub issues, not in this doc. |

49 | `DECISION\_SNAPSHOT\_REGRESSION.md` | **DRAFT ? decisions LOCKED** (D1?D11 locked 2026-06-25, #383). Zero code shipped for P1?P7. | No | Start P1+P2 (schema + stitch-plan factor); de-risk the P10 main.py split |

50 | `DESKTOP\_APP\_PLAN.md` | DRAFT (owner-gated, nothing built) | No | P0 compute-feasibility spike; relates to truly-local profile |

51 | `OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md` | PLAN/roadmap (owner-gated) | No (authoritative roadmap) | Gen-0?Gen-2; the owned-model track |

52 | `PERSONALIZATION\_PLAN.md` | DESIGN ? owner-adopted; Phase-0 | No | Per-user models (DB v5 store landed) |

53 | `UI\_ALPHA\_EXECUTION\_PLAN.md` | ACTIVE | No | khelsutra UI alpha track |

54 | `I18N\_PLAN.md` | Design (P0); **in implementation** (adapted) | No | Executed in `khelsutra/docs/FRIENDLY\_AND\_MULTILINGUAL.md` (P1?site/, P3?web/); P2 backend stays here, owner-gated |

55 | `POST\_166...RISK\_ASSESSMENT.md` | Plan for owner review | No | Milestones/risk doc |

56

57 | ### 3b. Foundational / strategy references (LIVING ? archive only if **superseded**, not when "done")

58 | Doc | Status | Note |

59 | ---|---|---

60 | `PLATFORM\_ARCHITECTURE.md` | Strategy/research | The `ComputeBackend`/`StorageBackend` registries ? realized by COMPUTE_DECOUPLED_SERVING |

61 | `COMMERCIALIZATION.md` | ? LIVING | Strategy; carries WASB-C3 |

62 | `VIDEO\_LOCALITY\_MODEL.md` | THINKING DOC | L0 fully-local = the truly-local profile |

63 | `HOSTING\_PLAN.md` | PROPOSAL | **? partly overtaken by ADR-013 + the shipped site ? verify (finding #5)** |

64 | `UI\_PROPOSAL.md` | PROPOSAL (partly shipped) | UX half **in implementation** ? `khelsutra/docs/FRIENDLY\_AND\_MULTILINGUAL.md`; currency banner added (finding #5 ? for this doc) |

65 | `STORAGE\_SHARING\_MODEL.md` | DRAFT (co-design pending) | ? |

```

66 | `ANALYZERS_AND_RECONCILERS.md` | DESIGN (owner-gated) | ? |
67 | `ROADMAP.md` / `CODE_MAP.md` / `HOW_RALLY_DETECTION_WORKS.md` /
`KHELSUTRA_PRODUCT_OVERVIEW.md` | LIVING references | Keep current; not archive targets |
68
69 ### 3c. Reports & terminal records (archive candidates once their track wraps)
70 | Doc | Status | Note |
71 |---|---|---|
72 | `ALPHA_LAUNCH_REVIEW_2026-07-04.md` | **REPORT ? DELIVERED 2026-07-04** (A13) ? owner-
facing alpha launch review: GO/NO-GO recommendation (conditional GO for the current 4-email
allowlist, gated on khelsutra #154 copy fixes + the owner browser-CUJ checklist; NO-GO for wider
invites until A10/A11 + SA least-privilege + retention), DoD scorecard, operator runbook,
observability/triage incl. the 2026-07-04 live failure-mode battery, privacy/consent audit,
known limits + rollback | Awaiting the owner GO/NO-GO record on `ALPHA_LAUNCH_READINESS.md` §10;
archive with the alpha cluster once recorded |
73 | `REAL_FOOTAGE_VALIDATION_REPORT.md` | validation complete | Motivates R2; archive with
the R-cluster |
74 | `RALLY_DETECTION_QUALITY_REPORT.md` | report | Quality snapshot; part of R-cluster |
75 | `QUALITY_ITERATIONS.md` | LIVING empirical log | **Keep** ? the paper-backbone
ablation log ([[paper-coauthor-aim]]) |
76 | `CODE_AUDIT_AND_TEST_HARDENING` . `HARDENING_LOOP_HANDOFF` . `DEFERRED_HARDENING_PLAN`
| COMPLETED ? **finish-archived 2026-07-02** (top-level stubs removed; absorbed into the
umbrella's *Prior audits*) | ? Full copies under `docs/archives/*-COMPLETED-2026-06-14.md`;
inbound links repointed (finding #4 resolved) |
77
78 ### 3d. Conventions / guidelines / templates / runbooks (permanent ? do not archive)
79 `PROJECT_DOC_TEMPLATE.md` . `NEXT_STEPS.md` (? refreshed 2026-06-21) . `BACKLOG.md` .
`EVALUATION.md` . `ABLATION_LAYER.md` . `SETUP_NEW_MACHINE.md` . **`CREATING_REELS.md`** (? end-
user usage guide ? install ? process ? pick ? reel; LIGHT tier; added 2026-06-23) .
**`CONVERT_GCS_VIDEOS_TO_HIGHLIGHTS.md`** (? operator runbook ? `gs://` videos ? reels on L4
Cloud Run; verified e2e 2026-06-23) . `TRACKNET_WSL_SETUP.md` . `README.md` . *(2026-06-21
relocation: the data-layer guides moved ? `data_pipeline/`; the field-ops + Roora docs ?
`archives/{field-pmf,roora-vendor}/`.)*
80
81 ### 3e. Already archived (`docs/archives/`)
82 `past_projects/DECOUPLED_COMPUTE/` (? 2026-06-21, #270) . `past_projects/low-regret-
moves/` . **`field-pmf/`** (field-ops + customer-feedback, 2026-06-21) . **`roora-vendor/`**
(annotation-vendor engagement, 2026-06-21) . `PER_VIDEO_OUTPUT_LAYOUT-SUPERSEDED-2026-06-21.md`
. `CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md` . `DEFERRED_HARDENING_PLAN-
COMPLETED-2026-06-14.md` . `HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md` .
`COMMERCIAL_READINESS_REVIEW.md` . `MONETIZATION_AUDIT.md` . `TEST_RUN_SUMMARY.md` . `research/`
. `checkpoints/` . `quality-iterations/`
. `past_projects/CODE_CLEANUP_AUDIT_2026-06-24-SUPERSEDED-2026-07-02.md` (absorbed ? umbrella) .
`past_projects/KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN-FOLDED-2026-07-02.md` (folded ? umbrella).
83
84 ### 3f. ADR log ? `docs/archives/decisions/DECISIONS.md`
85 **ADR-001 ? ADR-014** (14, all ratified). The compute/serving lineage: **ADR-005 ? 008 ?
009 ? 010 ? 011 ? 012 ? 013 ? 014** . ADR-014 was **RATIFIED 2026-06-21** (owner M0 sign-off for
COMPUTE_DECOUPLED_SERVING).
86
87 ## 4. Open due-diligence findings / actions (2026-06-21 review)
88
89 > **? Doc-restructuring done (2026-06-21, owner-directed):** field-ops/customer-feedback
? `archives/field-pmf/`, Roora annotation-vendor ? `archives/roora-vendor/`, and a no-ML **data
layer** carved into `docs/data_pipeline/`. Inbound + outbound links rewired; three dir READMEs
added (the `data_pipeline` one sets the "no segmentation/stitching" boundary).
90 1. **? RESOLVED (2026-06-21) ? `PER_VIDEO_WORKSPACE.md` supersedes
`PER_VIDEO_OUTPUT_LAYOUT.md`.** Folded forward the `video_slug` naming option (workspace §8.5) +
the `output/GX010125_run/` manual-prototype precedent; predecessor archived ?
`docs/archives/PER_VIDEO_OUTPUT_LAYOUT-SUPERSEDED-2026-06-21.md`; `CODE_MAP.md` repointed.
91 2. **? RESOLVED (2026-06-21) ? `NEXT_STEPS.md` refreshed** with a current 2026-06-21
state + a **cross-track prioritized pending list** ; the 2026-06-13 owned-model/quality body is
kept (clearly marked) as that track's detail.
92 3. **R-series / quality cluster is ACTIVE** (`R1`,`R2`,`RALLY_DETECTION_GAP_CLOSURE`,`RA
LLY_QUALITY_RESEARCH`,`QUALITY_ITERATIONS`,`REAL_FOOTAGE_VALIDATION_REPORT`,`RALLY_DETECTION_QUA
LITY_REPORT`). **Do NOT archive yet** ? archive as a *cluster* once R2 is promoted-or-rejected

```

as the gate + \$7.1's single number is attributed. (Owner already confirmed "leave as is", 2026-06-21.)

93 4. **RESOLVED (2026-07-02)** ? finish-archived the hardening stubs.** The three top-level pointer-stubs (`CODE\_AUDIT\_AND\_TEST\_HARDENING`, `HARDENING\_LOOP\_HANDOFF`, `DEFERRED\_HARDENING\_PLAN`) were removed; full archived copies remain under `docs/archives/`, inbound links (README, PLATFORM_ARCHITECTURE, DESKTOP_APP_PLAN) were repointed, and the sweep is summarized in the umbrella's *Prior audits (absorbed)*.

94 5. **HOSTING_PLAN.md** / **UI_PROPOSAL.md** are PROPOSAL but partly overtaken** by ADR-013 (DO dropped) + the shipped khelsutra site. **UI_PROPOSAL.md** banner added (2026-06-21)** ? its UX half is now tracked in `khelsutra/docs/FRIENDLY\_AND\_MULTILINGUAL.md`; the **I18N_PLAN.md** localization track is likewise in implementation there. **HOSTING_PLAN.md** banner still pending** (ADR-013 already noted in it).

95 6. **RESOLVED (2026-06-21)** ? **COMPUTE_DECOUPLED_SERVING** M0 approved + M1?M4 SHIPPED** (#279/#280/#282/#283); all \$8 gating Qs locked (D7?D17); M5+ default-OFF code batch authorized (owner residual: counsel ToS/DPA, real GCP prices, OAuth app, real cloud run within the D17 \$100/?10k cap).

96

97 ## 5. Maintenance

98 Update a row when its doc changes lifecycle; run \$2 before any archive; re-review \$4 at each clean-slate. This register + the \$2 checklist are the operational form of the **CLAUDE.md** archive policy and [[working-agreement-tracked-project-docs]].

99

9. assistant (2026-07-08T05:45:32.806Z)

Now let me read the ADRs. Let me find where 012/013/014 live in DECISIONS.md.

10. user (2026-07-08T05:45:34.452Z)

598:## ADR-012: Compute = GPU-native (L4/T4); TPU deferred; Cloud Run + GPU (scale-to-zero) is the supported serving model

661:## ADR-013: Drop DigitalOcean for compute ? GCP is the SOLE compute stack (training + serving)

686:## ADR-014: Compute-decoupled rally serving ? one pure core, profile-selected compute + storage

11. user (2026-07-08T05:45:35.746Z)

714 C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/docs/archives/decisions/DECISIONS.md

12. user (2026-07-08T05:45:41.206Z)

598 ## ADR-012: Compute = GPU-native (L4/T4); TPU deferred; Cloud Run + GPU (scale-to-zero) is the supported serving model

599

600 **Date:** 2026-06-20 . **Status:** RATIFIED (owner, in-session). **Refines ADR-011** (concretizes the "GCP = serving" path).

601

602 ### Decision

603

604 1. **Stay GPU-native (NVIDIA L4/T4) for all compute** ? local box, GPU VM, and the serving path ?

605 for the foreseeable roadmap. **Do NOT invest in TPU now.** Revisit TPU **only** if/when a large

606 ***owned* transformer (Qwen3-VL Gen-2+, owner-gated)** at ***sustained batch* scale** becomes the dominant

607 cost, and **only after a benchmark** proves it ? never speculatively.

608 2. **Cloud Run + GPU (L4), scale-to-zero, is the supported SERVING deployment model** for the

609 pay-per-video orchestration economics ("users pay us to orchestrate on our GCS bucket + VMs /

610 Cloud Run"): pay-per-request GPU-seconds = the exact per-video bill, zero idle-GPU cost, co-located

```

611     with `gs://khelsutra-rally-corpus`. This is the concrete form of ADR-011's deferred
612     **M3/M4**.
613     3. **The named optimization lever is GPU-accelerated video decode (NVDEC)**, not more
614     GPU FLOPs ? the
615     measured bottleneck is video *decode*, not the model.
616
617     ### Rationale
618     - **Workload isn't TPU-shaped.** WASB is a tiny (5.8 MB) PyTorch CNN doing per-frame
619     heatmap
620     inference; the owned model is a VLM. TPUs win on large, dense, *batched* matmul at
621     training scale;
622     streaming per-frame CNN inference with small/dynamic batches is a poor fit.
623     - **The real bottleneck is video decode, which TPU does nothing for.** Measured on the
624     2026-06-20 GCP
625     L4 M2 run: cv2 PNG extraction at ~12 fps with the **GPU idle during extraction**. The
626     lever that
627     helps is **NVDEC** (hardware decode on L4/T4) + parallel decode ? i.e. lean *into* the
628     GPU.
629     - **Portability / single-runtime.** TPU needs XLA (torch_xla / JAX) ? significant
630     fragile NRE to port
631     WASB (torch 1.11) + the pipeline, a second runtime to maintain, and **GCP-only** ? it
632     would break
633     the local ? VM ? Cloud Run parity the DECOUPLED storage/compute seams exist to
634     provide. One portable
635     PyTorch+CUDA image that runs identically everywhere is worth more than peak FLOPs
636     here.
637     - **Economics.** Cloud Run+GPU scale-to-zero meters per request = clean per-video
638     billing; TPU pricing
639     (v5e pods / queued / committed-use) targets sustained training and its regional
640     availability is
641     *narrower* than GPU (the M2 run already hit a Mumbai GPU stockout ? landed in
642     Singapore; TPU would be
643     worse).
644
645     ### Consequences
646
647     - The **deferred M3 (inference serving endpoint + per-video billing) and M4 (scale-to-
648     zero + cost
649     guard)** from ADR-011 become a **new tracked subproject ? "Cloud Run + GPU serving"**
650     ? to be spun
651     up (as a `docs/PROJECT_DOC_TEMPLATE.md`-based doc) **AFTER `docs/DECOUPLED_COMPUTE/`
652     is archived**
653     (owner sequencing). Seed scope below. **? REALIZED in ADR-014 (2026-06-21):** the
654     subproject is
655     `docs/COMPUTE_DECOUPLED_SERVING/`, **renamed** to reflect that **truly-local is a co-
656     equal profile**
657     (not just cloud serving); ADR-014 expands this 7-point seed into the M0?M10 plan.
658     - **First subproject work item = containerize `python -m backend.worker`** for Cloud
659     Run. This is the
660     known hard part: WASB's torch-1.11 conda env makes the image non-trivial (two runtimes
661     ? the
662     indexer venv + the `wasb` conda env). The decoupled storage/compute seams + the native
663     runner
664     (`detector_impl="native"`, PR #256) make the *code* portable; packaging is the gap.
665     - Keep the **storage/compute decoupling seams** intact so the SAME container runs local
666     / GPU VM /
667     Cloud Run (the per-video workspace convention, in progress, reinforces this).
668     - **Edge auth is a hard prerequisite** before any public serving (engine API is CORS
669     ``, no auth) ?
670     carried from the appliance track.
671     - **Region/capacity reality:** GPU stockouts are common (Mumbai was fully stocked-out on
672     2026-06-20);
673     scale-to-zero + multi-region/flexible-zone selection is part of why Cloud Run+GPU
674     fits.

```

650 - TPU is explicitly ****out of scope****; this ADR records ***why*** so it isn't relitigated.

651

652 ##### Seed scope for the "Cloud Run + GPU serving" subproject (formalize post-archival)

653 1. Containerize `backend.worker` (indexer venv + `wasb` conda env; NVDEC-capable base image).

654 2. Cloud Run + GPU (L4) service: `infer` over a `gcs:///` workspace; scale-to-zero.

655 3. Per-video cost ledger (on the `jobs` table) = the per-customer bill (M3).

656 4. Cost guard + concurrency caps (M4).

657 5. Edge auth (tenancy boundary) before public exposure.

658 6. NVDEC GPU-decode to kill the extraction bottleneck.

659 7. DPA / consent gate for non-owner uploads (productization, not compute).

660

661 ## ADR-013: Drop DigitalOcean for compute ? GCP is the SOLE compute stack (training + serving)

662

663 ****Date:**** 2026-06-20 . ****Status:**** RATIFIED (owner, in-session: "drop DO and document it? record it and bake it for all future work"). ****Supersedes the DO=training split in ADR-011; refines ADR-010 + ADR-012.****

664

665 ### Decision

666

667 1. ****DigitalOcean is DROPPED as a compute provider.**** ****GCP is the sole compute stack**** for BOTH training and serving. (ADR-011's "DO=training, GCP=serving" split is retired ? GCP does both. The DECOUPLED scope decision in ADR-011 ? M0?M2 + M3.5 ? still stands.)

668 2. ****Standing default for ALL future compute work = GCP**** (a GPU VM today via `deploy/gcp/`; Cloud Run+GPU for serving per ADR-012). Do ****not**** re-evaluate DO/Paperspace for compute without a new ADR.

669 3. The `\$200` DO credits, if used at all, go ****only to non-compute**** that DO bills cleanly (CPU droplets / a site host / Spaces storage / bandwidth) ? ****never GPU****.

670

671 ### Rationale (researched 2026-06-20, against DO's own docs)

672

673 - ****DO GPU is neither enabled on the account nor a clean free-credits path.**** The DO API `v2/sizes` returns ****0 GPU sizes**** (GPU Droplets not enabled); GPU Droplets are ****North-America-only**** (NYC2/TOR1/ATL1/RIC1/AMS3 ? no India/Singapore); and DO's own docs state that for ****newer customers DO charges the team's primary payment method**** when GPU Droplets are created (anti-abuse) ? so GPU spend hits the ****card****, not the promo credit. The ADR-011 premise ("200 credits fund DO training") does ****not**** hold.

674 - ****Paperspace (now DO Gradient) free GPU is subscription-gated**** (a paid "G"/Growth plan; high-end A100/H100 needs the \$39/mo Growth plan), and Paperspace bills ****separately**** from DO Droplet credits ? the \$200 credit doesn't fund it either.

675 - ****GCP works today:**** quota granted (`GPUS\_ALL\_REGIONS=1`), bucket + SA + `\$100` budget provisioned, and M1/M2/M3.5 ran on an L4 for ~\$12/run with no subscription. Real money, but trivial per run ? and ****one**** portable PyTorch+CUDA stack (local ? GPU VM ? Cloud Run), no second provider to maintain.

676

677 ### Consequences

678

679 - ****Compute provisioning = `deploy/gcp/` only.**** `SETUP\_NEW\_MACHINE.md` repointed to GCP; `deploy/digitalocean/README.md` marked ****DEPRECATED****; `HOSTING\_PLAN.md`'s "All-DigitalOcean" compute option retired.

680 - ****`deploy/digitalocean/` scripts are KEPT**** ? `bootstrap.sh` / `mount\_scratch.sh` / `gpu\_setup.sh` are ****provider-agnostic**** Linux-GPU setup the GCP runbook reuses; only the DO-specific ****provisioning**** (doctl droplet create, Spaces) is dropped. (Optional later cleanup: move them to `deploy/common/`.)

681 - The shared DO CPU droplet ****`claude-do-rally-test`**** (another session's testbed) is ****owner-coordinated**** cleanup, separate from this decision; the freshly-rotated ****DO API token**** can be retired once that droplet is gone.

682 - ****Storage stays bucket-abstracted**** (`gcs:///` primary). The `s3` backend still speaks any S3-compatible store, so ****DO Spaces remains a *possible* storage target**** even though DO ****compute**** is dropped.

683 - ****Capacity reality:**** GCP GPU stockouts are common (asia-south1/Mumbai was fully stocked-out 2026-06-20 ? the M2 run landed an L4 in asia-southeast1/Singapore; weights still land in the asia-south1 bucket). Sweep zones/regions; scale-to-zero (ADR-012) is part of why Cloud Run+GPU fits.

```

684 - **DECOUPLED_COMPUTE** archives on M2 completion with a note that the D0 leg was
**dropped** (this ADR), not "delivered on D0".
685
686 ## ADR-014: Compute-decoupled rally serving ? one pure core, profile-selected compute +
storage
687
688 **Date:** 2026-06-21 . **Status:** ? **RATIFIED** (owner, 2026-06-21 ? M0 design sign-off;
build M1+). **Refines ADR-011 + ADR-012** (it *realizes* their deferred M3/M4; it does **not**
supersede them); **builds on ADR-013** (GCP sole compute stack). Full design:
`docs/COMPUTE_DECOUPLED_SERVING/`.
689
690 ### Evolution / lineage (trace the idea)
691 `ADR-005` (permissive licensing ? engine + image stay MIT/Apache/BSD) . `ADR-008`
(train==serve featurizer single-sourcing ? the parity discipline this extends to local?cloud) .
`ADR-009` (**local-first delivery** ? origin of the truly-local profile) . `ADR-010` (DECOUPLED
DO?GCP) . `ADR-011` (**deferred M3 serving + M4 billing**; overrode "rent nothing / not a
service"; $06 gates = DPA/consent, collector `md5`, WASB-C3) . `ADR-012` (**Cloud Run+GPU =
serving model**; NVDEC lever) . `ADR-013` (**GCP sole compute stack**) . **? ADR-014** (this).
Non-ADR: realizes `PLATFORM_ARCHITECTURE.md`'s `ComputeBackend`/`StorageBackend` registries; the
truly-local profile = `VIDEO_LOCALITY_MODEL.md` L0 + `DESKTOP_APP_PLAN.md` `LocalDesktop`;
informed by the archived `DEPLOY_REPORT_GCP_2026-06-20` (cross-GPU sub-pixel parity finding).
692
693 ### Decision
694
695 1. The rally engine is a **PURE** function of (input bytes + config) ? outputs**, in three
deterministic stages ? **DETECT** (`DetectorRunner.run_predict`?CSV), **WINDOW**
(`trajectory_to_action_windows`, CPU-pure), **STITCH** (`VideoCompiler.compile_highlights`,
ffmpeg) ? that are **NEVER** branched on the deployment profile**.
696 2. A **`ComputeBackend` / deployment profile** decides **WHERE** the core runs (local
process vs a Cloud Run GPU job); a **`StorageBackend`** decides **WHERE** I/O comes from** (local
FS / USB / mounted Google Drive / bucket). **Both are selected by ONE startup/setup config**
(`deployment.profile` + `compute_target`), realizing the PLATFORM_ARCHITECTURE
`local|hosted|byo_worker` registry.
697 3. Ship at least two profiles: **truly-local** (local GPU + **in-process stitch**, no
cloud) and **cloud-serving** (a Cloud Run GPU job runs detect **AND** stitch), plus **cpu-mock**
for CI and a reserved, **deferred** `byo_worker` hybrid.
698 4. In cloud-serving, **STITCH** runs on Cloud Run** (co-located with detect) so its
CPU/wall-time + egress are **metered** into a per-job cost ledger alongside `gpu_active_seconds`
? stitching is compute-intensive and its cost must be accounted for. (Truly-local stitch is
intentionally **un**metered.)
699 5. **Parity** is enforced at **WINDOW + STITCH** (byte-exact for a fixed trajectory);
**DETECT** is byte-exact only **same-GPU** ? cross-GPU is sub-pixel-different (deploy report
2026-06-20), so it is judged at the **rally/window level**, never CSV bytes.
700 6. Land **smallest-safe-first, one PR per milestone (M0?M10)** , **DEFAULT-OFF** where a
change touches core callers (e.g. the configurable output base), each guarded by golden-fixtures
+ a **profile-parity test** ; default flips ship a **Launch Report** .
701
702 ### Rationale
703
704 - The seams already exist: `_resolve_runner` is the single compute seam, the
`StorageBackend` ABC + 4 backends + the worker fetch/put are built, and the stub detector
(`replay_csv`) gives a **$0 deterministic** core. Realizing a profile system is **composition,
not a rewrite** .
705 - Separating **WHERE** (ComputeBackend, coarse, per-job) from **HOW** (DeviceContext,
fine, per-GPU) from **WHAT-bytes** (StorageBackend) keeps each axis independently testable and
avoids the main trap ? conflating them.
706 - **Metering stitch on Cloud Run** is the only way the owner's real serving cost (GPU +
CPU re-encode + egress) is attributable per job; stitch-on-laptop hides it.
707 - Aligns with ADR-012 (GPU-native) + ADR-013 (GCP sole stack) + the 2026-06-20 deploy
report; the truly-local profile preserves the ADR-009 privacy/offline self-host story.
708
709 ### Consequences
710
711 - **Positive:** one core, provably identical across modes; CI proves the decoupling for
**$0** (stub + storage fakes + a profile-parity test); new compute targets (byo_worker,

```


SkyPilot, Modal) slot into the enum without core changes; ****cost becomes visible per job**** (incl. stitch).

712 - ****Cost/new surface:**** a new `deployment`+`compute\_target` config surface, a Cloud Run ****dispatch shim****, a ****containerized GPU+ffmpeg worker image****, a ****cost-ledger migration****, and ****edge auth**** must all be built + maintained. Cross-GPU detection is ****not**** byte-identical ? parity asserted at the window level. The `byo\_worker` tier carries a real ops/support burden (deferred, power-user only).

713 - ****Risk-managed:**** every core-touching change is default-OFF + golden/parity-gated + Launch-Reported; real-GPU runs are ****owner-side**** with a budget guard and a posterity DEPLOY_REPORT (`docs/COMPUTE\_DECOUPLED\_SERVING/runlogs/`).

714 - ****Owner gates ? the 5 gating ones LOCKED 2026-06-21**** (project doc §2 ****D7?D12****): D7 cold-start = ****scale-to-zero****; D8 pricebook = ****versioned DB table, USD-internal/INR-display****; D9 edge-auth = ****Cloudflare Access**** (verify Cf-Access JWT, lock ingress); D10 = ****free-tier data-for-reels trade**** (adults-only, derived-data, counsel-gated to M9); D11 quality-floor = ****Gemini-ON until the owned model clears the floor**** + D12 the ****data-flywheel harness**** (capture?store?eval?goldenize). Plus ****3 more locked at M0 sign-off (2026-06-21, §2 D13?D15):**** ****D13**** profile is user-switchable anytime (advertised feature; equivalent output, not bit-identical cross-GPU); ****D14 NORTHSTAR**** = mobile/cloud-native point-at-Google-Drive ? reel back in Drive next to the source (implies a new ****gdrive-api` OAuth\*\*** backend = M12; resolves the gdrive-in-cloud question); **\*\*D15\*\*** auto-detect SUGGESTS, the user CONFIRMS ? execution (esp. cloud spend) is never a surprise. Still open (non-gating): stitch co-located vs split; NVDEC vs libx264; the `<2GB` cap; the dir name.

13. user (2026-07-08T05:45:54.102Z)

docs\CODE_MAP.md:22:Top-level tree and the ****placement rule**** for each kind of file. When unsure, prefer an existing

docs\CODE_MAP.md:162:[Omitted long matching line]

docs\archives\CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md:156:- Shallow section replacement behavior.

docs\archives\CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md:361: -
`test\_config\_shallow\_section\_replacement`: verifies loader's shallow **{**defaults, **file_data}** (whole section from file replaces defaults section; no deep merge of sub-keys).

docs\archives\HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md:150: - Restored the strong original `test\_get\_returns\_dict\_for\_nested\_sections\_legacy\_compat` (exact name + all legacy dict/nested/scalar/leaf-key/ default/fallback guards that the weaker shim replacement had dropped).

docs\HOW_RALLY_DETECTION_WORKS.md:32:From the per-frame `(x, y)` we compute the shuttle's ****velocity**** (displacement between frames).

docs\COMMERCIALIZATION.md:63:| ****torchvision**** detectors (Faster R-CNN/RetinaNet) | BSD-3 | No |
? LIVE | The chosen replacement; person = COCO class (see
`backend/pipeline/detectors/person\_detector.py` `\_PERSON\_CLASS\_ID`). |

docs\OWNED_MODEL_IMPLEMENTATION_PLAN.md:52: wholesale replacement. *(Open: the VLM-ceiling number is open-domain; our amateur footage may differ ? re-measure.)*

docs\archives\roora-vendor\ROORA_LABELING_GUIDELINES.md:192:****Boundary placement is the non-negotiable part:****

docs\archives\roora-vendor\ROORA_LABELING_GUIDELINES.md:342:- `forced\_error` = the opponent's shot CAUSED the miss (pace, placement, deception; the player barely

docs\COMPUTE_DECOUPLED_SERVING\README.md:90:### 4.4 Compute placement ? and ***why** stitch runs on Cloud Run* `[design]`

docs\COMPUTE_DECOUPLED_SERVING\README.md:140:2. ? ****Stitch placement ? co-located**** in the detect job (= D4; one metered unit; the job table supports a future CPU-split non-breaking, revisit only on bursty-compile metrics).

docs\KHELSUTRA_PRODUCT_OVERVIEW.md:161:- It is ****not a replacement for a coach**** ? it's a tool that saves coaches time and adds insight.

docs\POST_166_WORK_ITEMS_MILESTONES_AND_RISK_ASSESSMENT.md:7:****Key outcomes from #166 (merged c558372):**** Hardening loop (T1-T5 + finale #165) COMPLETE and archived. NEXT_STEPS reframed with explicit P3 start signal (harness LOVO litmus for person features) + full prioritized candidate list. Gate A wired (opt-in, default-OFF safety per served regression litmus). Loop scheduler (019ec6d07da5) deleted; no replacement created. Handoff and archived audit updated terminal. Master also includes #167 (packaging).

docs\POST_166_WORK_ITEMS_MILESTONES_AND_RISK_ASSESSMENT.md:21:[Omitted long matching line]

docs\POST_166_WORK_ITEMS_MILESTONES_AND_RISK_ASSESSMENT.md:45:[Omitted long matching line]

docs\R2_EVAL_METRIC_DESIGN.md:156:- ****Default-additive ? ablation-gated replacement:**** R2 is reported next to tIoU (augment). The ***gate***

docs\R2_EVAL_METRIC_DESIGN.md:188:- **Gate: AUGMENT now, then ablation-gate the eventual replacement, then deprecate-with-docs ? ADOPTED.**

docs\PER_VIDEO_WORKSPACE.md:6:[Omitted long matching line]

docs\ROADMAP.md:32:pipeline objectively and (2) train an offline replacement for the paid Gemini

docs\RALLY_DETECTION_QUALITY_REPORT.md:288:aggressively (so the gate trims **real** rallies instead). The honest, served replacement is **W2**

docs\STORAGE_SHARING_MODEL.md:185:13. On-device milestone placement: academy pilot boxes first, consumer app after the WSL-free

docs\archives\past_projects\AUDIT_REPORT_khelsutra-guru_2026-07-02-COMPLETED-2026-07-02.md:216:[Omitted long matching line]

docs\archives\past_projects\AUDIT_REPORT_khelsutra-guru_2026-07-02-COMPLETED-2026-07-02.md:369:[Omitted long matching line]

docs\TRACKNET_WSL_SETUP.md:67: "velocity_threshold": 0.02 // normalised shuttle displacement/sec to count as "in flight"

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md:135:[Omitted long matching line]

docs\archives\past_projects\CODE_CLEANUP_AUDIT_2026-06-24-SUPERSEDED-2026-07-02.md:46:| **D5** |

Follow the CODE_MAP placement rules. No new top-level files without a placement-rule match. Tools go in `tools/`, standalone run-drivers in `scripts/`, eval library code in `backend/eval/`. | `docs/CODE\_MAP.md` \$0 | Keeps the repo navigable as it evolves. |

docs\archives\past_projects\CODE_CLEANUP_AUDIT_2026-06-24-SUPERSEDED-2026-07-02.md:266:-

CODE_MAP.md `docs/CODE\_MAP.md` ? placement rules, gotchas, and the cross-cutting footguns list.

docs\archives\past_projects\DECOUPLED_COMPUTE\07-COMPUTE-ABSTRACTION.md:92: def to(self, model_or_tensor): ... # placement (cuda/mps/cpu)

14. user (2026-07-08T05:45:55.212Z)

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md:135:[Omitted long matching line]

15. user (2026-07-08T05:46:05.410Z)

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_alpha-alliv1_2026-07-06.md:64:Rough (M8 ledger not landed): Cloud Run L4 `gpu\_active` ? 1,800 + 2,412 + 956 + 1,302 ? **6,470 s** ? 1.8 L4-hours**, plus the diag VM ? 1 L4-hour (g2-standard-4). At ~US\$0.85 / L4-hour on-demand ? US\$2.4 for the whole diagnostic session. The proxy path is ~50 % fewer detection GPU-seconds per clip going forward.**

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_alpha-alliv1_2026-07-06.md:66:## 8. Teardown / final state (\$0 audit)

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_alpha-alliv1_2026-07-06.md:68:- Diag VM **deleted** (never stopped): `gcloud compute instances list` ? **0 items**; `gcloud compute disks list` ? **0 items** (no orphan boot disk). ? \$0.

docs\COMPUTE_DECOUPLED_SERVING\runlogs\README.md:26:8. **Teardown / final state** ? for cloud, the `\$0` audit (`gcloud compute instances list` + `disks list` empty); **DELETE**, never just stop ([gcp-vm-always-delete]).

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_cp-rebuild_2026-07-05.md:62: after hashing. Cost of the encounter here: one L4 re-detection (~US\$0.4), harvested as the

docs\COMPUTE_DECOUPLED_SERVING\runlogs\RUNLOG_truly-local_TEMPLATE.md:41:## 4. Parity sanity (optional, \$0)

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md:100:- 2 L4 worker execs (Run A 157.7 s + Run B 109.6 s ? 268 s of L4 time) ? **~\$0.5?0.7**. Off-box A/B: **\$0** (local). No VM, no image build. GCS read for the 3 downloaded clips (~5.1 GiB) is intra-project.

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_validation-latency_2026-07-07.md:137:**Cost (heuristic):** one L4 detect exec (`rally-worker-64bhj`, ~222 s) ? **\$0.15?0.55** (rate-dependent, dominant) + control-plane CPU for validation+stitch (~960 s @ 2 vCPU) ? **~\$0.05** + reel egress (581 MB, if downloaded to internet) ? **~\$0.07** + storage/ops **< \$0.02** + Gemini **\$0** (skip_ai_handoff). **\$0.3?0.7 total**, ~one L4 exec as projected ? but the stitch's real cost is **concurrency** (~16 min of control-plane heavy CPU), not dollars.

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_2026-07-04.md:3:> **Date:** 2026-07-04 · **Owner-authorized** GCP spend (D17, \$300 envelope) · deleted every

docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_2026-07-04.md:4:> resource after ([gcp-vm-always-delete]); \$5 residual audit = \$0 of ours).

```
docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_2026-07-04.md:52:- Builds: 4
× ~$0.16 (cu113 base reuse, cu118, batch) ? $0.5
docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_2026-07-04.md:53:- Runs:
baseline + cu118 + batch64 + batch128(OOM, short) + probes ? ~$2?3 of L4 time
docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_2026-07-04.md:54:- Idle:
**$0** (scale-to-zero; GPU bills only during execution)
docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_2026-07-04.md:55:- **Total
well under the $300 / D17 envelope.**
docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_2026-07-04.md:57:## 5.
Teardown ? residual audit ($0 of ours)
docs\COMPUTE_DECOUPLED_SERVING\runlogs\DEPLOY_REPORT_cloud-serving_2026-07-04.md:70:No
GPU/compute of ours remains; idle cost = $0.
docs\COMPUTE_DECOUPLED_SERVING\runlogs\RUNLOG_cloud-serving_exec-status-
fastfail_2026-07-07.md:70:### A ? classifier vs live history (read-only, $0)
docs\COMPUTE_DECOUPLED_SERVING\runlogs\RUNLOG_cloud-serving_exec-status-
fastfail_2026-07-07.md:132:- **A:** $0 (read-only `list_executions` + `get_execution`).
docs\COMPUTE_DECOUPLED_SERVING\runlogs\RUNLOG_cloud-serving_exec-status-
fastfail_2026-07-07.md:136:- Total: well under $1.
```

16. user (2026-07-08T05:46:07.116Z)

```
01-EXISTING-PLANS-AND-GAP.md
02-COMPUTE-MAP-AND-CONTRACT.md
03-PLATFORM-OPTIONS.md
04-BUILD-AND-TEST-PLAN.md
05-COST-ATTRIBUTION.md
06-RISKS-AND-OPEN-QUESTIONS.md
07-COMPUTE-ABSTRACTION.md
DEPLOY_REPORT_GCP_2026-06-20.md
HANDOFF.md
README.md
```

17. user (2026-07-08T05:46:15.975Z)

```
1      # 05 · Cost attribution ? per-video inference & per-run training billing
2
3      **Status:** PROPOSAL (net-new design) · **Date:** 2026-06-20
4
5      > **Side goal:** a non-owner user runs a job and gets the **exact billing for that
video** (inference), and
6      > training cost can be **segmented similarly** (per run, optionally amortized per
contributed video). This
7      > is wholly new ? the `jobs` table has **no cost columns** today
([`database.py:188`](../../backend/infrastructure/database.py)).
8      > The good news: a job is already the unit of work, and serverless platforms already
meter per invocation ?
9      > so per-video billing is mostly *bookkeeping on top of seams that exist.*
10
11     ## 1. The core idea
12
13     **A job is the billing unit.** An `infer` job is **1:1 with a video** (keyed by
`video_id`/MD5), so the
14     job's cost *is* the per-video bill. A `train` job is **1:1 with a training run**, so its
cost is the run
15     cost ? which we then optionally **split across the videos that fed it**.
16
17     ```
18     cost_usd(job) = ?_resource usage(job, resource) × rate(resource, pricebook_version)
19     ```
20
21     Resources to meter: **GPU-seconds** (the big one), CPU-seconds, wall-seconds, **egress
GB** (reel
22     downloads ? the line `HOSTING_PLAN.md` flags), storage GB-hours, and **external-API
calls** (Gemini
23     tokens/calls ? the provider already returns metrics: `analyze_video` ? (segments,
```

```

metrics)`).
24
25     ## 2. Metering mechanics (cheapest accurate source first)
26
27     | Source of GPU-seconds | When | Accuracy |
28     |---|---|---|
29     | **Serverless per-invocation report** (RunPod Serverless / Cloud Run / Modal) | M3
inference endpoint | **Best** ? the platform bills exact GPU-seconds per job; that number *is*
the bill. **This is the strongest argument for a serverless `infer` topology**
([03](03-PLATFORM-OPTIONS.md)). |
30     | **Worker-measured `gpu_active_seconds`** (timestamp around the GPU section, × the
known `gpu_type` rate) | DO Droplet / VM where there's no serverless meter | Good ? separate
`gpu_active_seconds` from `wall_seconds` so you don't bill the user for the bucket-sync/upload
time. |
31     | **Provider billing API reconciliation** (monthly) | all | Truth source ? reconcile the
sum of job costs against the invoice; carry a `reconciliation_factor`. |
32
33     Egress is counted from bytes downloaded (R2 = zero egress ? simplest). Gemini cost comes
from the metrics
34     the provider already returns.
35
36     ## 3. Schema delta (extends the existing `jobs` table; no new datastore)
37
38     Add via the existing `_add_column_idempotent` / `PRAGMA user_version` migration ladder:
39
40     ```sql
41     ALTER TABLE jobs ADD COLUMN owner_id          TEXT;      -- who pays (the non-owner user /
tenant)
42     ALTER TABLE jobs ADD COLUMN compute_backend TEXT;      -- runpod_serverless |
gcp_cloudrun | do_droplet ?
43     ALTER TABLE jobs ADD COLUMN gpu_type          TEXT;      -- L4 | A4000 | T4 ?
44     ALTER TABLE jobs ADD COLUMN gpu_active_seconds REAL;
45     ALTER TABLE jobs ADD COLUMN wall_seconds      REAL;
46     ALTER TABLE jobs ADD COLUMN bytes_out         INTEGER;   -- egress
47     ALTER TABLE jobs ADD COLUMN gemini_cost_usd   REAL;
48     ALTER TABLE jobs ADD COLUMN pricebook_version TEXT;
49     ALTER TABLE jobs ADD COLUMN cost_usd          REAL;      -- computed total
50     ALTER TABLE jobs ADD COLUMN cost_breakdown_json TEXT;    -- {gpu: .., egress: .., api:
..}
51
52     -- price book (versioned; rates change, and differ per provider/gpu_type)
53     CREATE TABLE IF NOT EXISTS pricebook (
54         version TEXT, resource TEXT, gpu_type TEXT, unit_price_usd REAL, currency TEXT,
effective_at TIMESTAMP
55     );
56
57     -- training-run cost attribution across contributing videos
58     CREATE TABLE IF NOT EXISTS training_contributions (
59         training_job_id TEXT, video_id TEXT, weight REAL      -- weight = clip-seconds share (or
1/N)
60     );
61     ```
62
63     The **`pricebook` is seeded from the same provider-grounded GPU set as the device
profiles**
64     ([07](07-COMPUTE-ABSTRACTION.md) §5) ? DO/Paperspace now (A4000 / L40S / A100), GCP
later (T4 / L4 / A100) ?
65     so **one `{provider, gpu_type}` table carries both the $/hr rate (here) and the device-
tuning knobs (there)** ,
66     keyed off the `gpu_type` that `run_telemetry`'s `nvidia-smi` probe already yields.
67
68     `owner_id` is the per-tenant scoping the multi-tenant model
([`PLATFORM_ARCHITECTURE.md`](../PLATFORM_ARCHITECTURE.md) §4)
69     will need anyway ? adding it now keeps "a non-owner user is billed for their jobs" a
first-class property.

```

```

70
71     ## 4. Per-video inference billing
72
73     1. An `infer` job's `cost_usd` + `cost_breakdown_json` **is** the per-video bill (job is
1:1 with the video).
74     2. Expose `GET /api/jobs/{id}/cost` and a per-video cost history (`GET
/api/videos/{video_id}/cost`), scoped
75         by `owner_id`.
76     3. Optional resale markup: a `margin` config ? `price_to_user = cost_usd × (1 +
margin)`; store both so the
77         user sees a transparent breakdown and you keep your margin auditable.
78
79     **Example breakdown (illustrative):** 3-min L4 inference + a 40 MB reel download + one
Gemini call ?
80     `gpu: 180 s × $0.00022/s = $0.040` + `egress: 0.04 GB × $0 (R2) = $0` + `gemini: ~$0.05`
? **~$0.09/video**.
81
82     ## 5. Per-training cost segmentation
83
84     1. A `train` job's `cost_usd` is the **run cost** (e.g. Gen-0 ? a few cents; a future
VLM run would be $-tens).
85     2. To "segment similarly," attribute the run cost across the N videos/labels in that run
via
86         `training_contributions` ? **equal split (1/N)** or **weighted by clip-seconds
contributed**. This
87         yields **"cost per video contributed to training," the symmetric analog of the per-
video inference bill.
88     3. Record `training_job_id` on the produced weight bundle so a served model's lineage
carries its training
89         cost (useful for amortizing training into a per-inference price later, if you
choose).
90
91     ## 6. Trust & reconciliation
92
93     - **Transparency:** always show the user the `cost_breakdown_json` +
`pricebook_version`, not just a total.
94     - **Reconciliation:** metered cost ? provider invoice but not exact (overhead, rounding,
idle). Reconcile
95         monthly: `reconciliation_factor = invoice_total / ? metered_cost`; surface drift if it
exceeds a few %.
96     - **Spend guard:** a per-`owner_id` rate limit + a global hard cap (the same cap that
protects the $200)
97         means a leaked API key can't run up someone else's bill or drain the budget.
98
99     ## 7. Why this is cheap to build
100
101     It rides entirely on seams that exist: the **`jobs` table** (already the unit of work +
has `compute_backend`
102     in the design), the **per-video output layout**, the **`tenant_id`** the platform doc
already plans, and the
103     **ComputeBackend** (which knows `gpu_type` + rate). The only genuinely new pieces are
the **pricebook**, the
104     **usage capture in the worker**, and the **attribution table** ? all small, and none
requiring a new
105     datastore.
106

```

18. user (2026-07-08T05:46:26.294Z)

```

backend\api\job_worker.py:201:    (= {gpu_seconds, wall_seconds, egress_gb, gemini_usd,
gpu_type}) × the versioned pricebook, and
backend\api\job_worker.py:205:    ``billing.enabled=True`` AND a real calibrated pricebook,
every recorded cost is **0** until a
backend\api\job_worker.py:206:    FOLLOW-UP wires the worker to emit usage
(gpu/wall/egress/gemini). M8 lands the ledger + pricebook

```

```

backend\api\job_worker.py:208:    With the placeholder pricebook the cost is 0 regardless.
backend\api\job_worker.py:221:        pricebook_version=billing_cfg.pricebook_version,
backend\api\job_worker.py:228:        "Recorded cost for job %s (pricebook=%s).",
backend\api\job_worker.py:230:        billing_cfg.pricebook_version,
backend\billing.py:3:``cost_usd(job) = ?_resource usage(job, resource) × rate(resource,
pricebook_version)``
backend\billing.py:6:(the Bangalore market). The pricebook is a versioned DB table
(``backend/infrastructure/database.py``);
backend\billing.py:7:this module is the pure arithmetic over a resolved ``{resource:
unit_price_usd}`` rates dict, so it
backend\billing.py:10:**Default-OFF:** with the seeded ``placeholder-v0`` pricebook (every rate
``0.0``) every cost is
backend\billing.py:11:``0.0`` ? nothing user-facing changes until the owner inserts a real,
calibrated pricebook version.
backend\billing.py:18:# Metered resource keys ? shared by the usage dict, the pricebook rows,
and the breakdown.
backend\billing.py:28:# Resources whose cost = amount × unit_price (GEMINI_USD is excluded ? it
is already a USD amount).
backend\billing.py:38:    key to its USD unit price (already resolved for this pricebook version
+ gpu_type by the DB
backend\billing.py:40:    Missing amounts/rates default to ``0.0`` ? with the placeholder
pricebook the total is ``0.0``.
backend\config\models.py:491:    ``pricebook`` DB table and **displayed in INR** at
``fx_inr_per_usd``. The default pricebook
backend\config\models.py:492:    version is the seeded ``placeholder-v0`` (every rate ``0.0`` ?
cost ``0``) until the owner
backend\config\models.py:493:    inserts a real, calibrated version and points
``pricebook_version`` at it.
backend\config\models.py:495:    ? Even with ``enabled=True`` + a real pricebook, recorded cost
stays **0** until a follow-up
backend\config\models.py:497:    pricebook + the recording seam; usage emission (the metering
hooks) is a separate step.
backend\config\models.py:500:    pricebook_version: str = (
backend\config\models.py:501:        "placeholder-v0" # the seeded $0 version; owner flips to a
real one
backend\main.py:981:    null until billing records it (default-OFF / the placeholder pricebook ?
0)."""
backend\infrastructure\database_jobs.py:316:    def get_pricebook_rates(

```

[Showing results with pagination = limit: 25]

19. user (2026-07-08T05:46:27.266Z)

```

1      # DEPLOY REPORT ? cloud-serving, GPU-resident NVDEC decode rollout (#506/#507),
2026-07-07
2
3      > Durable record of shipping the **GPU-resident WASB decode**
(`decode_backend=nvdec_resident`, issue #506 / PR #507) to the live `cloud-serving` path:
rebuild + cut over the `rally-worker` L4 Cloud Run Job image (torch 1.11 ? **2.6.0+cu124 +
PyNvVideoCodec 2.1.0 + baked NVDEC libs**), smoke the cpu + nvdec paths on the real L4, then
**enable nvdec via a reversible Job env var** and canary it. Project `gen-lang-
client-0306026826`, region `asia-southeast1`. Honest-limits in $6.
4
5      ## 1. TL;DR
6
7      | Stage | Result |
8      |---|---|
9      | Merge | ? PR #507 squash-merged to master (`44a2f47`); full local gate + CI green |
10     | Build image | ? `rally-worker:latest` ? `sha256:465a9250` (Cloud Build 12m7s) |
11     | Job cutover | ? `jobs update --image ? :latest` (prior digest `sha256:3bfefaf6` =
rollback) |
12     | Phase A ? cpu smoke (torch 2.6) | ? 180s clip: `torch 2.6.0+cu124 cuda True`, 17
windows, WASB 201.9s |
13     | Phase B ? nvdec smoke (fast) | ? 180s clip: 19 windows, WASB **143.7s (~1.4× vs
cpu)**, 2934 vis ? cpu 2920 |

```

```

14 | Phase B ? nvdec smoke (**hard GX010128**) | ? **26479/45298 visible ? VM 26481** (?
0.004%), 85 windows, WASB 1138s (~40 fps) |
15 | **Enable** | ? `WASB_DECODE_BACKEND=nvdec_resident` set on the Job (reversible) |
16 | Canary (env-driven) | ? 180s clip, **no** decode `--set` ? env var selected nvdec: 19
windows, 2934 vis, WASB 151.4s (== explicit-`--set` nvdec) |
17
18 **Bottom line:** the GPU-resident decode path runs on the real Cloud Run L4 and
**reproduces the validated VM trajectory** (26479 ? 26481 visible on GX010128), so the golden-
gate F1 (0.574 ? baseline 0.582 @ SERVED) transfers to serving. nvdec is enabled via a Job env
var with an instant rollback.
19
20 ## 2. Resources
21 - **Worker:** `rally-worker` Cloud Run **Job**, image `asia-
southeast1-docker.pkg.dev/gen-lang-client-0306026826/rally/rally-worker:latest`
(`sha256:465a9250?`), 8 vCPU / 32 GiB / 1x NVIDIA L4, task-timeout 3600 s, GCSFuse mount of
`gs://khehsutra-rally-serving`, `WASB_WEIGHTS=/gcs/models/wasb_badminton_best.pth.tar`,
**`WASB_DECODE_BACKEND=nvdec_resident`** (the enablement lever).
22 - **Image change (PR #507):** WASB conda env py3.8+torch1.11 ? **py3.10 + torch
2.6.0+cu124 + PyNvVideoCodec==2.1.0**; `libnvcuvid`/`libnvidia-encode` (from driver 580.159.03)
baked in a multi-stage layer. cpu decode path byte-unchanged; nvdec is a new, cache-isolated
path.
23 - **No VM used** for this deploy (the image build is Cloud Build; smokes run on the
existing Job).
24
25 ## 3. Exact commands
26 ``bash
27 # Build + push (from master @ 44a2f47)
28 gcloud builds submit --config=deploy/cloudrun/cloudbuild.yaml --project gen-lang-
client-0306026826 .
29 # Cut the Job over to the new image (re-resolves :latest to the new digest)
30 gcloud run jobs update rally-worker --region asia-southeast1 --image ?/rally/rally-
worker:latest
31 # Smoke (direct Job execute; skip_ai = secret-free; --set decode_backend per leg). ^@@^
arg delimiter.
32 gcloud run jobs execute rally-worker --region asia-southeast1 --async \
33 --args='^@@^infer@@-video-uri=gs://khehsutra-rally-corpus/videos/<clip>@@-out-
uri=gs://khehsutra-rally-serving/rel506/<label>@@-segmenter=wasb_hybrid@@-
set=deployment.compute_target=local_gpu@@-set=indexing.detector_impl=native@@-
set=indexing.skip_ai_handoff=true@@-set=indexing.wasb.batch_size=64@@-
set=indexing.wasb.prefetch_batches=4@@-set=indexing.wasb.timeout_sec=0@@-
set=indexing.wasb.decode_backend=<cpu|nvdec_resident>\'
34 # ENABLE (reversible): the native runner reads WASB_DECODE_BACKEND first, so this flips
all dispatches
35 gcloud run jobs update rally-worker --region asia-southeast1 --update-env-vars
WASB_DECODE_BACKEND=nvdec_resident
36 # ROLLBACK options
37 gcloud run jobs update rally-worker --region asia-southeast1 --remove-env-vars
WASB_DECODE_BACKEND # nvdec ? cpu (same image)
38 gcloud run jobs update rally-worker --region asia-southeast1 --image ?/rally/rally-
worker@sha256:3bfefaf6? # full revert to the torch-1.11 image
39 ``
40
41 ## 4. Run results (executions)
42
43 | Exec | Clip | decode | via | visible / frames | windows | WASB s | status |
44 |---|---|---|---|---|---|---|---|
45 | rally-worker-p2wd6 | mahadevpura_smoke_180s | cpu | `--set` | 2920 / 5395 | 17 | 201.9
| ? |
46 | rally-worker-clkwc | mahadevpura_smoke_180s | nvdec_resident | `--set` | 2934 / 5395 |
19 | 143.7 | ? |
47 | rally-worker-6htsw | **GX010128** | nvdec_resident | `--set` | **26479 / 45298** | 85
| 1138.0 | ? |
48 | rally-worker-cnzxw | mahadevpura_smoke_180s | nvdec (via **env**) | Job env var | 2934
/ 5395 | 19 | 151.4 | ? |
49

```

50 - GX010128 nvdec on Cloud Run: ****26479 visible vs the validated VM run's 26481**** (? 2 frames / 45298 = 0.004%) ? the serving trajectory matches the golden-gate trajectory, so its ****F1 0.574 @ SERVED (? baseline 0.582)**** holds in prod.

51 - Fast-clip nvdec vs cpu (same tuning batch 64 / prefetch 4): WASB ****143.7 s vs 201.9 s (~1.4x)****, near-identical trajectory (? visible 0.3%).

52

53 **## 5. Rollout state + rollback**

54 - ****Enabled**** on `rally-worker` via `WASB\_DECODE\_BACKEND=nvdec\_resident`. All future dispatches use the GPU-resident path.

55 - ****Rollback (instant, same image):**** `--remove-env-vars WASB\_DECODE\_BACKEND` ? back to the (torch-2.6) cpu path. ****Full revert:**** `--image sha256:3bfefaf6?` ? the prior torch-1.11 image.

56 - ****Durability (landed):**** PR #508 set `indexing.wasb.decode\_backend: nvdec\_resident` in the `cloud-serving` preset. But the worker Job has ****no `DEPLOYMENT\_PROFILE`****, so it never applies that preset ? a follow-up (`chore/506-release-hardening`) makes the control plane ****forward**** the resolved `indexing.wasb.{decode\_backend,batch\_size,prefetch\_batches}` to the worker (`orchestration.py`), so the preset is the true source of truth on the dispatch path. Until the control plane is redeployed with that fix, the `WASB\_DECODE\_BACKEND` env var above is what keeps nvdec live (and remains a valid manual override afterwards). ? The env var also revealed that the preset's `batch\_size=64/prefetch\_batches=4` were ****never reaching the worker**** pre-fix (dispatch forwarded only `skip\_ai\_handoff`) ? the forwarding fix closes that latent gap too.

57 - ****Update (later 2026-07-07) ? forwarding is now LIVE.**** Rebuilt the image from master `66756ad` (= #507+#508+#509) ? ****`sha256:09d2edf7`**** and redeployed BOTH surfaces to it: control-plane service rev ****`rally-control-plane-00017-qbd`**** (100% traffic, booted healthy ? uvicorn up, cloud-serving profile ACTIVE, edge-auth on) and the `rally-worker` Job. So the dispatch forwarding + the preset default now govern the worker from the image; the `WASB\_DECODE\_BACKEND` env var is retained as a redundant belt-and-suspenders override (both say nvdec). ? The FIRST control-plane deploy attempt used a ****stale image**** `sha256:465a9250?` (built from `44a2f47`, pre-#508/#509 ? no forwarding); it became rev 00016 and was superseded by 00017. Not live-dispatch-verified end-to-end (the `/api/process` path is behind Cloudflare Access ? no CF JWT this session); forwarding is deployed + unit-tested. ****Control-plane rollback:**** `gcloud run services update-traffic rally-control-plane --region asia-southeast1 --to-revisions rally-control-plane-00015-kp9=100` (last pre-nvdec revision).

58

59 **## 6. Honest limits / not-verified**

60 - Smokes ran with `skip\_ai\_handoff=true` (secret-free) so the persisted "rallies" are the local candidate windows, not Gemini-refined ? this validates the ****decode + detection + windowing**** path, not the AI handoff (unchanged by this PR).

61 - F1 parity is asserted by the ****offline golden gate**** (10 clips @ SERVED) + the ****trajectory reproduction**** above (Cloud Run 26479 ? VM 26481 visible on GX010128); no on-Cloud-Run re-scoring against labels was run (the trajectory match makes it redundant).

62 - The control-plane ****service**** was redeployed to `sha256:09d2edf7` (see §5 Update), so the API path now runs the merged code (forwarding + persisted `decode\_backend` metadata). The worker Job is the release-critical surface (it's the one that decodes) and was validated first.

63 - NVDEC driver-ABI: the baked `libnvcuvid 580.159.03` works on the current Cloud Run L4 host driver (proven by the successful decode). A future Cloud Run host-driver change could in principle skew ? re-smoke on any base-image/driver bump (the Dockerfile comment flags this).

64

20. user (2026-07-08T05:46:37.849Z)

1 # DEPLOY REPORT ? control-plane durable rebuild (CP-rebuild P9), 2026-07-05

2

3 > The §4.8/§9 execution record for `docs/CONTROL\_PLANE\_DURABLE\_REBUILD.md` (spec #487, owner-confirmed).

4 > Everything here ran live against project `gen-lang-client-0306026826`, region `asia-southeast1`,

5 > on 2026-07-05 (14:10?14:5x UTC). Honest-limits in §6: what this run does NOT prove.

6

7 **## 1. What shipped in the image**

8

9 Master `257f469` ? the full P-row sweep, all merged on verified LGTMs the same day:

10 P1 #489 (generator pins + 8Gi) · P2 #490 (GCS-streamed upload staging) · P4 #492 (submit-time


```

11 MD5 + cache-reuse) · P5a #491 (poll budget 3900 s + capabilities caps) · P6 #493 (cloud-
poll
12 heartbeat) · P3 #494 (restart-safe recovery + read-path re-hydrate) · P8 #495 (SPA re-
vendedored
13 at khelsutra `90d1d90` = #153/#155/#158/#162 baked; platform-stable `ui_version` hash).
14 khelsutra side: #162 (60-min SPA ceilings) on khelsutra master `90d1d90`.
15
16 ## 2. Build ? pin ? deploy ($4.8 steps 1?4)
17
18 | Step | Command / object | Result |
19 |-----|-----|-----|
20 | Image build | `gcloud builds submit --config=deploy/cloudrun/cloudbuild.yaml` from a
clean worktree at `257f469` | build `47877373` **SUCCESS in 7m33s** ? digest
`sha256:3bfefaf6ad4a?7671b964` (moved from the prior live `31475f27?`) |
21 | Worker re-pin | `gcloud run jobs update rally-worker --image ?/rally-worker:latest` |
job resolves the new digest |
22 | CP deploy (auth ON) | `gcloud run services replace service-auth-on.yaml` (generated by
`deploy/cloudrun/gen_service_yaml.py --edge-auth on` from the clean clone) | rev **`rally-
control-plane-00012-wxk`** · **8Gi** / 2 vCPU · imageDigest == `3bfefaf6?` |
23 | Pins in the baked config `[verified pre-deploy]` | base64-decoded the manifest's
wrapper payload | `indexing.default_segmenter="wasb_hybrid"` + `security.max_part_mb=24` +
`edge_auth_enabled=true` present ? **D5 works from a clean clone** |
24 | Upload-staging lifecycle | `gcloud storage buckets update --lifecycle-file ?` |
`videos/uploads/**` age-2d Delete rule live (bucket previously had NO lifecycle policy) |
25
26 ## 3. Posture ($9, on the final auth-ON revision 00012)
27
28 - `GET /api/health` (identity token) ? **200**
`{"status":"ok","sport":"badminton","default_segmenter":"wasb_hybrid"}` ? the segmenter pin is
visible in the health payload itself.
29 - `POST /api/process` without a CF JWT ? **401** (edge-auth ON).
30 - `https://api.khelsutra.guru/api/health` ? **302** (Cloudflare Access intercept).
31
32 ## 4. Verification battery (short auth-OFF window, rev `00013-rlc` ? the precedented
A1-e2e pattern)
33
34 The auth-off deploy replaced the instance ? **fresh, empty SQLite: exactly the
2026-07-05
35 post-restart state** the spec's DoD demands be replayed.
36
37 | Proof | What ran | Result |
38 |-----|-----|-----|
39 | **Restart replay (#471/#485 ? THE loss shape)** | `POST /api/query
{"video_id":"c2d45319?f55369"}` against the wiped DB | **29 play_segments returned** ? re-
hydrated from `outputs/<md5>/` on demand (`metadata.source=cloud_run` rows). The July-5 morning
failure shape now recovers by itself |
40 | **Streamed upload staging (P2/#480)** | 2×768 KiB parts ? `/api/upload/init ? part/0 ?
part/1 ? complete` | `path=gs://?/videos/uploads/7383?p9-stream-proof.mp4`,
`video_id=dc65ef78?`; **server-side object `md5Hash` (hex) == the streamed video_id EXACTLY**,
`component_count` empty (non-composite) ? browser-staged uploads stay md5-addressable, so the
instant cache path holds for the tester flow |
41 | **Capabilities (P5a/#480 item 5)** | `GET /api/capabilities` | `max_upload_mb=4096`,
`max_part_mb=24`, `default_segmenter=wasb_hybrid`, `deployment.cloud_available=true` |
42 | **Heartbeat (P6/#482 CP slice)** | `GET /api/jobs/{id}` during the live detection |
`progress = "segmenting (wasb_hybrid) on cloud GPU ? 5 min elapsed"` ? later ticks 6, 7 ? the
silent 22-minute spinner is dead |
43 | **Jobs queue (#481)** | `GET /api/jobs?status=all` | lists the running job with its
live progress |
44 | **Fresh detection on the new image** | the same submit dispatched `rally-worker-2mz8t`
(see $5 for WHY it dispatched) | job ? **done, 48 segments**, `compute.path=cloud_run`, exec
`?/rally-worker-2mz8t`, **wall 14m04s** (14:23:41?14:37:45Z). Clarifier: the bucket's
`MahadevpuraSingles.MP4` hashes to `9154737?` ? a DIFFERENT file than the morning's 2.68 GB
browser upload (`c2d45319?/29 rallies); 48 matches the A1-e2e record for this original, and its
`outputs/` were idempotently rewritten |
45 | **Cache proof ? the literal $9 box (#484)** | immediate re-submit of the same `gs://`

```

```

path (same instance, DB now warm) | done cached=true, path=cached, 48 segments` in 11 s;
`latestCreatedExecution` unchanged (`2mz8t`) ? zero GPU on the re-submit |
46
47     ## 5. Findings (honest ? two real ones)
48
49     1. Composite originals bypass the submit-time cache (documented P4 boundary, now
measured).
50     `videos/MahadevpuraSingles.MP4` was bucket-uploaded as a composite object
51     (`component_count=4`) ? GCS exposes no `md5Hash` ? `submit_time_video_id` returns
None ?
52     the instant-done path can't fire. By itself that would only degrade to the drain-tier
DB
53     check; combined with (2) it re-dispatched.
54     2. Fresh-DB + composite-original ? the bucket cache is never consulted on the drain
path.
55     Cloud Run replaced the rev-00013 instance between the replay query and the submit
56     (`MANUAL_OR_CUSTOMER_MIN_INSTANCE` churn in the logs) ? the drain hashed on an empty
DB,
57     found no video row, and dispatched ? even though `outputs/<md5>/report.json` existed.
The
58     drain checks only the DB after hashing; the bucket tier lives on the submit path
(md5-known)
59     and the read path. Residual gap, bounded: it needs BOTH a composite/md5-less
source AND
60     an instance swap; browser-staged uploads (the tester flow) always carry `md5Hash` and
take
61     the instant path. Follow-up filed on #484: consult `cached_process_result` on the
drain
62     after hashing. Cost of the encounter here: one L4 re-detection (~US$0.4), harvested
as the
63     fresh-detection proof below.
64
65     ## 6. What this run does NOT prove (stays on the owner's $1.1 CUJ)
66
67     - The 2.5 GB browser upload through Cloudflare Access on the real domain (the $9
headline
68     memory-survival proof under load) ? the synthetic upload proves the staging PATH, not
the
69     4 GiB-scale memory profile. The owner's official 5-minute CUJ covers it on this build.
70     - SPA-rendered UX (progress copy, medium picker, pickable rallies) ? API-level proofs
only here;
71     the vendored snapshot provenance (`ui_version.spa_commit=90d1d90`, content-hash guard
in CI)
72     is what ties the served UI to the verified khelsutra master.
73     - Concurrent-duplicate submit collapse (#481) was regression-proven in CI (12-thread
test) but
74     not re-driven live in this window.
75
76     ## 7. Restore + final posture
77
78     `services replace service-auth-on.yaml` ? rev `rally-control-plane-00014-7hk` (edge-
auth ON,
79     8Gi, same digest `3bfefaf6?`). Re-verified: health (identity token) 200, `POST
/api/process`
80     sans CF JWT 401. The `p9-stream-proof.mp4` staging object was deleted (the age-2d
lifecycle
81     rule would have collected it regardless). The auth-off battery jobs' DB rows died with
the
82     window's instance ? by design and now harmless: their results live in the bucket and any
read
83     re-hydrates on demand ($4 row 1 is exactly that proof).
84
85     ## 8. Tracker updates riding this PR
86
87     - `docs/CONTROL_PLANE_DURABLE_REBUILD.md` $7: P9 ? (this runlog); status header

```

```

note.
88 - `docs/ALPHA_LAUNCH_READINESS.md`: DoD item 1 annotated ? the infra blocker behind the
failed
89 2026-07-05 owner CUJ is deployed-fixed (rev 00012-wxk); the official CUJ pass remains
the
90 owner's checkbox. A1 row cross-references this runlog.
91 - Issues: #480 + #471 + #485 closed with live evidence; #484 kept open (composite-gap
follow-up);
92 #482 stays open (worker-side heartbeat + per-job runlog).
93

```

21. user (2026-07-08T05:46:51.355Z)

```

1 # Control-Plane Durable Rebuild ? kill the upload OOM, ground state in the bucket,
unblock DoD #1
2
3 > **Status:** `IN PROGRESS` ? **P9 EXECUTED 2026-07-05**: live = rev **00012-wxk** (8Gi,
digest
4 > `3bfefaf6?`, edge-auth ON, pins live). $9 residual = the owner's $1.1 browser CUJ
(three boxes);
5 > flip to `DONE` when it passes . **Owner:** avidullu . **Created:** 2026-07-05 .
6 > **Last updated:** 2026-07-05
7 > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE)
8 > **Tracking anchors:** $7 progress tracker is the source of truth; indexed in
`docs/README.md`;
9 > pointer in memory `current-state` (+ `docs/NEXT_STEPS.md` once work starts post-
confirmation).
10 > **Relation to existing docs:** executes the serving hardening that
11 > [ALPHA_LAUNCH_REVIEW_2026-07-04.md](ALPHA_LAUNCH_REVIEW_2026-07-04.md) $1.1 (the
5-minute browser
12 > CUJ ? tracker **DoD #1**) is blocked on; child of
13 > [COMPUTE_DECOUPLED_SERVING/](COMPUTE_DECOUPLED_SERVING/README.md) (ADR-014);
consolidates issues
14 > #471 #480 #482 #483 #484 #485 (plus the already-shipped #479-config, #481, and
khelsutra
15 > #153/#154(#155)/#157(#158)).
16 > **Honesty note:** each claim is tagged `[verified]` (checked against code or live
logs),
17 > `[researched]` (external platform fact ? re-verify during execution), or `[design]`
(proposed).
18
19 ---
20
21 ## 0. TL;DR
22
23 The first real owner-driven browser CUJ (2026-07-05) proved the product works end-to-end
? 29 rallies
24 detected, pickable in the UI ? **when the control plane survives**. It also proved the
deployed CP
25 does not survive real usage: a 2.68 GB browser upload staged into RAM-backed `/tmp` on a
**4 GiB**
26 instance ? OOM ? restart ? ephemeral SQLite wiped ? the 29 rallies (already written to
the bucket)
27 were never ingested ? empty UI, unrecoverable by reload `[verified]`. This spec
consolidates the fixes
28 into **one image rebuild + redeploy**: ?8 GiB CP with upload staging **streamed to the
bucket** (never
29 assembled in tmpfs) . **bucket-grounded durable state** (the DB becomes a cache that re-
hydrates from
30 `outputs/<md5>/`) . **submit-time MD5** (makes re-runs instant AND restarts recoverable
? it is the
31 key that re-associates an orphaned job with its output) . longer, honest poll ceilings .
the rev-00011
32 config pins **baked into the committed manifest generator** so a regenerate can never

```

silently drop

33 them · the current SPA (#153/#154/#158) re-vendored. The deploy step (**P9**) is owner-gated on

34 confirming this spec; the code rows (P1?P8) are each one small PR.

35

36 ## 1. Problem & goal

37

38 **What happened (2026-07-05, all diagnosed live ? full record in memory

39 [[alpha-cuj-testing-2026-07-05]]:**

40

41 1. 95 MB upload parts were 413'd by the Google Front End ? the service has no HTTP/2 annotation, so

42 ingress is HTTP/1 with a ~32 MiB per-request body cap (#480). Config rev 00011 set

43 `security.max\_part\_mb=24` and small uploads started working `[verified]`.

44 2. Cloud jobs ran `yolo\_hybrid`, not the validated `wasb\_hybrid` (#479) ? the engine default

45 (`backend/config/models.py:181`) was never pinned. Config rev 00011 pinned it

46 `[verified]`.

47 3. **The killer:** a 2.68 GB upload assembled into the CP's RAM-backed `/tmp` on a 4 GiB instance ?

48 OOM ? restart 08:14Z ? the ephemeral SQLite **and the in-memory done-marker poll** were wiped

49 (#471). The worker finished (`done rc=0` 08:59:51Z) and wrote 29 rallies to the bucket, but no CP

50 was listening ? never ingested ? the UI showed nothing and reload could not recover it

51 `[verified]`. The reel was recovered manually from `outputs/<md5>/intervals.csv`.

52 4. Timeouts are dishonest for real footage: the SPA gives up at 10 min

53 (`web/src/hooks/useIngestJob.ts:18`, `web/src/api/client.ts:181` in the khelsutra repo) and the CP

54 poll budget is 30 min (`backend/infrastructure/execution\_policy.py:54`), while a 6-min 1080p60 clip

55 takes WASB ~22 min on the L4 `[verified]`.

56 5. No MD5 reuse: the same clip was detected **3x** that day ? three paid L4 runs for one video

57 (#484) `[verified]`.

58 6. The gs:// re-run CUJ **PASSED in the UI** (29 rallies pickable) because nothing staged into tmpfs ?

59 the product works when the CP survives `[verified ? owner screenshot]`.

60

61 **Goal.** A tester can, from a browser at `https://api.khelsutra.guru` (edge-auth ON), upload a real

22. user (2026-07-08T05:46:52.836Z)

19:[Omitted long matching line]

20:[Omitted long matching line]

36:[Omitted long matching line]

23. user (2026-07-08T05:46:59.939Z)

1 - RALLY_QUALITY_RESEARCH.md ? **? Foundation doc** for closing the local-vs-Gemini rally-detection gap: the measured gap + the code-grounded over-merge diagnosis ("shuttle moving ? rally"), an inventory of the eval/experiment harness + its gaps, the strengthened evaluation framework (segmental Edit/F1@k that make over-segmentation first-class, golden-corpus + active learning, Gemini-agreement as a *shadow* signal), an

adversarially-verified, cited external-research synthesis (TAL/TAS boundary heads, racket-sports rally-state cues, small-data, eval rigor), a **prioritized cheap?durable roadmap with falsifiable success criteria**, and (\$7) a **tracked project plan** ? work-item checklist, sequencing, the single quality **number** attributed to the effort, and a **definition of done** (project completes when all tiers land + the number is recorded). The plan that the QUALITY_ITERATIONS.md log executes against; ties together EXPERIMENT_HARNESS / MULTI_SIGNAL_FUSION_PLAN / ANALYZERS_AND_RECONCILERS / OWNED_MODEL_IMPLEMENTATION_PLAN.

2 - RALLY_DETECTION_GAP_CLOSURE.md ? **Tracked execution

doc (2026-06-18)** for closing the remaining rally-detection accuracy gap **per video, before launch**, driven by the growing golden corpus. First ground-truth clip (GX010141) reframed it: boundaries are tight (R2/R3/R4 holds) ? the gap is **detection**: (G1) recall (local & Gemini both miss ~half on hard clips, and *different* rallies ? complementary), (G2) precision (local over-fires), (G3) **rally-END over-extension on the quick-return-after-point** (shuttle still moving ? needs a *player-state* signal). Levers (owner-gated): player-kinematics + shuttle-in-hand / sustained-invisibility end-rule with reverse-timeline backtrack, `rally\_gate` Gate A wiring, complementary fusion. **Peer** of `RALLY\_QUALITY\_RESEARCH.md` (execution tracker, not research); §7 is the source of truth.

3 - PROMINENT_COURT_DETECTION.md ? **Design + experiment (2026-06-19, default-OFF, owner-gated flip)** for the **multicourt G2 over-fire**: identify the **prominent (foreground) court** (the one covering the majority of the frame) and keep only rallies whose shuttle lives inside it ? re-defining the shuttle signal `(x,y,visible)` ? `(x,y,visible-AND-inside-prominent-court)` *before* windowing, so a background/adjacent-court shuttle never forms a rally. 2D floor polygon (MVP) ? **pseudo-3D play-volume** (refinement, so high clears aren't clipped). Born from the GX010142 rally-#1 background-court FP; reuses `fusion\_features.point\_in\_polygon` + `FusionConfig.court\_polygon`. §8 progress tracker is the source of truth.

4 - RALLY_DETECTION_QUALITY_REPORT.md ? **Technical-paper checkpoint (2026-06-17)** of rally-detection quality: the honest measured state (over-segmentation milestone tolerance-F1 0.09170.323 at n=11, p=0.002; merge-rate 0.69470.150; first real-footage ground-truth at-par with START at Gemini parity), the task-faithful R2 metric, a detailed Gemini comparison, next steps + expected wins, areas yet to be explored, and a publishability / path-to-submission verdict. Renders to a polished PDF; the externally-shareable "toward a paper" summary of `RALLY\_QUALITY\_RESEARCH.md` + `QUALITY\_ITERATIONS.md`.

5 - EVALUATION.md ? How we measure and improve quality.

6 - EXPERIMENT_HARNESS.md ? Design (P1 SHIPPED) for A/B + shadow + SxS experimentation on rally-detection variants with **zero production regression** ? experiments are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift, rollback = flip a flag). Composes the existing `backend/eval/` machinery (`calibration`, `regression`, `metrics`); implemented in `backend/eval/experiment.py`.

7 - QUALITY_ITERATIONS.md ? **Living log** of rally-detection quality iterations (hypothesis ? what shipped ? measured/expected effect ? lesson). The reverse-chronological record; terminal-state snapshots get archived under `archives/quality-iterations/`.

8 - ANALYZERS_AND_RECONCILERS.md ? Design (owner-gated, NOT started): a **signal-fusion framework** for rally detection ? producers emit confidence-scored **signals** at three temporal **scopes** (frame detectors / window analyzers / session analyzers), a learned **scorer** weights them, and an auditable **report-first reconciliation** pass lints the decision against the evidence. The spine: **no special-casing in the decision path** (signals?scorer, never `if/else`). Worked examples: a service-fault analyzer, a warm-up/practice span detector, and consuming a per-frame shuttle-state detector. A peer to `EXPERIMENT\_HARNESS.md` and `MULTI\_SIGNAL\_FUSION\_PLAN.md` (ADR-007 modular rally layer).

9 - PERSONALIZATION_PLAN.md ? Design (owner-adopted **hybrid**;
Phase-0 landing) for the **per-user personalization feature on a generalization-first baseline**: a thin per-user *tuning* layer gated by the honest small-N statistics (a **min-N promotion floor** + a **structural-guard exemption**, never a service gate) + a **contribution flywheel** (free-tier videos pool into the owner-trained baseline; **a tool that gets better the more you help it**). Records the **locked owner decisions** (identity, storage, `min\_n` semantics, Gate-A default, free-tier pooling). Built so far: per-user promotion gate (#141) + held-out-USER learning curve (#142).

10

11 - GOLDEN_REGRESSION_FIXTURES.md ? **Design (owner-gated, NOT started)**: video-free, non-attributable regression fixtures so a clean clone runs the rally-seg suite with **no video/GPU/DB**. Two layers ? the post-detector freeze + **dual gates** (characterization + quality-floor) reusing the `backend/eval/` stack, and the **SEALED-FIXTURES** privacy filter (de-attribution-at-source, tiered public/key-gated, shuttle-only public tier). Carries a cited **IN/US/EU legal memo** (derived-data IP/privacy), a threat model, and a **§7 progress tracker + definition-of-done**. Extends [GOLDEN_DATA_SHARING.md](data_pipeline/GOLDEN_DATA_SHARING.md).

12

13 ## ? Documentation status & governance

14 - [AUDIT_REPORT_khelsutra-guru_2026-07-02 (? COMPLETE / ARCHIVED)](archives/past_projects/AUDIT_REPORT_khelsutra-guru_2026-07-02-COMPLETED-2026-07-02.md) ? the five-repo `khelsutra-guru` workspace code audit +

its folded-in **Remediation tracker** (P0?P28 all closed). **Archived 2026-07-02** once every remediation wave landed and the last residuals were closed or filed as issues (#443 CORS, #444 corpus Phase-2, #445 GCS eval; + #301/#327/#332/#384). Absorbed the prior `CODE_CLEANUP_AUDIT_2026-06-24` + June-2026 hardening sweep (both archived).

15 - DOC_STATUS.md ? **living Documentation Status & Archival Register**: every foundational/tracked doc's lifecycle + archive-readiness, the **archival due-diligence checklist** (*honestly update ? then archive*), and the open doc-hygiene findings. Review here; update rows as docs complete.

16

17 ## Product / platform plans

18 - ALPHA_LAUNCH_READINESS.md ? **Active tracked project** (`DRAFT`, created 2026-07-03): **cross-repo alpha readiness tracker** after the golden corpus reached 15 videos. Separates the hosted product-alpha lane (shareable upload?process?pick?reel) from the owned-model/commercial rally-detection gate. A10/A11 are recorded only for a limited product alpha capped at fewer than 25 individually allowlisted testers; #172 and #486 remain the model/public/commercial gates. S7 is the source of truth for the concrete next PR sequence: corpus portability, n=15 baseline refresh, ablation/fusion reruns, hosted serving, alpha copy/launch review.

19 - CONTROL_PLANE_DURABLE_REBUILD.md ? **Active tracked project** (`DRAFT` owner-confirm-gated, created 2026-07-05): **the consolidated control-plane rebuild** that unblocks tracker **DoD #1** after the first owner browser CUJ (2026-07-05) proved the deployed CP fails real usage (2.68 GB upload ? 4 GiB RAM-`/tmp` OOM ? restart ? ephemeral SQLite wiped ? detected rallies never ingested). One image rebuild + redeploy: **8 GiB CP** · upload staging STREAMED to the bucket (never tmpfs) · bucket-grounded durable state (DB = cache, re-hydrate from `outputs/<md5>`) · submit-time MD5 (instant re-runs + restart recovery) · honest poll ceilings · rev-00011 config pins baked into `gen_service_yaml.py` · SPA re-vendored (#153/#154/#158); bundles #482 progress + #483 golden-gated downsample when ready. Consolidates #471 #480 #482?#485. S7 tracker is the source of truth; **P9** (deploy) is owner-gated on spec confirmation.

20 - [COMPUTE_DECOUPLED_SERVING/](COMPUTE_DECOUPLED_SERVING/README.md) ? **Active tracked project** (2026-06-21, `DRAFT` owner-gated): **the productionization successor** to DECOUPLED_COMPUTE (its deferred M3/M4). **One pure core** (DETECT?WINDOW?STITCH) that never branches on deployment profile, with a `ComputeBackend` (where it runs) + `StorageBackend` (what bytes) selected by **one startup config**. Ships **truly-local** (local/USB + local GPU, in-process stitch, zero cloud) + **cloud-serving** (upload<2GB / gdrive / bucket ? a **Cloud Run GPU** job runs detect AND stitch, metered into a per-job cost ledger) + **cpu-mock** (CI). Includes [TESTING_STRATEGY.md](COMPUTE_DECOUPLED_SERVING/TESTING_STRATEGY.md) (profile-PARITY proof for \$0), [architecture.svg](COMPUTE_DECOUPLED_SERVING/architecture.svg), a [runlogs/](COMPUTE_DECOUPLED_SERVING/runlogs/) posterity dir, and **ADR-014** ([lineage ADR-009?010?011?012?013?014](archives/decisions/DECISIONS.md)). S7 M0?M10 tracker is the source of truth.

21 - [DECOUPLED_COMPUTE/](archives/past_projects/DECOUPLED_COMPUTE/README.md) ? **DELIVERED + ARCHIVED** (2026-06-21). **Lifted the rally pipeline** onto a **rented cloud-native GPU** doing bucket-driven `train` + `infer`: storage/compute seams (PRs #246?#259) + the **M3.5 native WASB runner**; **M0?M2 + M3.5** proven on a GCP L4 (M2 ? `weights/0.1.0-l4/`). Delivered on **GCP** ? DigitalOcean dropped for compute ([ADR-013](archives/decisions/DECISIONS.md)). **M3/M4** (serving endpoint + per-video billing + scale-to-zero) deferred ? the "Cloud Run + GPU serving" subproject ([ADR-012](archives/decisions/DECISIONS.md)). Snapshot under `docs/archives/past_projects/DECOUPLED_COMPUTE/` (incl. `DEPLOY_REPORT_GCP_2026-06-20.md`).

22 - PACKAGING_PLAN.md ? **Active tracked project** (`IN PROGRESS`, 11 PRs merged 2026-06-21/22): **ship the engine** as a **downloadable batteries-included Python app** ? `pip install` ? `indexer --app` boots a local webserver and opens the **bundled KhelSutra SPA** single-origin, **torch-free** (LIGHT tier), cross-platform, no native component (decision D10/O1). Built ON the decoupled-compute seam. **Supersedes the DESKTOP_APP_PLAN.md** narrative. S7 tracker is source of truth (P0a?e ?, P1a/b ?, P2-launcher/P2a/P2d/P2e ?). **Next**: P3 GPU validation + screencast ? P2b-install ? P2c wizard (cross-repo) ? P4 (ONNX/train/annotate = north star). Memory `[[packaging-app-plan]]` / `[[next-session-packaging-app]]`.

23 - DESKTOP_APP_PLAN.md ? **Superseded narrative** (by `PACKAGING_PLAN.md`; D10 = self-hosted webserver, NOT a desktop app). **Original scoping doc** (2026-06-18, `design`): package the local pipeline as a **cross-platform desktop app** for non-technical players ? plug in the camera's USB/SD card, click once, get back **only the highlight reels**, fully on-device (no upload, no original-video bloat). The centerpiece is the **de-WSL native-compute port** (run WASB natively per-OS: Linux/Windows CUDA · macOS MPS/CPU ? drop WSL2), and its **make-or-break unknown**: is CPU acceptable on a typical **no-GPU**?

enthusiast laptop (WASB is ~3.4× real-time on a laptop 2070S)? \$7 tracker = **P0 compute-feasibility spike ? P1 native detector ? P2 app shell ? P3 packaging/installers**. Realizes the **L0 fully-local** tier of VIDEO_LOCALITY_MODEL.md; is the **`LocalDesktop` ComputeBackend** of PLATFORM_ARCHITECTURE.md §3.2; the smaller/export-clean owned model (OWNED_MODEL_IMPLEMENTATION_PLAN.md) is one answer to the compute crux.

24

25 ## Practical / ongoing

26 - ALPHA_TESTER_ONBOARDING_FAQ.md ? **? Tester-facing onboarding + FAQ for the hosted alpha** (send-ready; owner-approved copy per the A10/A11 guardrails ? workflow claims only). Sign-in via email OTP at `api.khelsutra.guru` · what footage works · the upload?detect?pick?reel flow with honest timing expectations from the 2026-07-05 CUJ · resumable links · privacy pointer (`khelsutra.guru/privacy`) + delete-on-request · ground rules (personal invite, rights to footage, non-commercial) · known limits (4 GB cap, conservative quality gate being tuned, watermark).

27 - CREATING_REELS.md ? **? End-user guide:** install the LIGHT-tier app ? `indexer --app` ? process a video ? pick rallies ? download a watermarked reel. Task-first usage walkthrough (web UI + the API chain), with a \$6 pointer to the cloud compute option. LIGHT tier only (offline `motion` / optional Gemini; no on-device GPU or one-click cloud). Install mechanics ? root [`README.md`](../README.md); packaging roadmap ? PACKAGING_PLAN.md.

28 - CONVERT_GCS_VIDEOS_TO_HIGHLIGHTS.md ? **? Operator runbook:** `gs://` bucket videos ? highlight reels on the approved **L4 Cloud Run** (scale-to-zero, \$0 idle). Copy-pasteable: build image ? create the GPU Job ? cloud-serving control plane ? process(cloud)?query?compile?download ? prove the path (`verify\_compute\_path.py`) ? \$0 teardown. **Verified e2e 2026-06-23.** Raw deploy knobs ? ../deploy/cloudrun/README.md.

29 - SETUP_NEW_MACHINE.md ? **? The canonical "spin up a box ? ready for training & inference" runbook.** Copy-pasteable, human-followable, **verified end-to-end on DigitalOcean (2026-06-20)** ? but **compute is now GCP-only (ADR-013)**; see ../deploy/gcp/README.md: credentials ? provision ? bootstrap ? suite (922 passed) ? secret-free CPU inference (no GPU/keys) ? real Gemini inference ? Gen-0 training, with pointers to the GPU/native-detector + bucket paths. The `deploy/digitalocean/\*` scripts are provider-agnostic and reused on GCP.

30 - DATA_IN_GCS.md ? **? The corpus source-of-truth map: where every piece of data lives in `gs://khelsutra-rally-corpus`** (canonical layout · current contents + drift · how `backend.worker` consumes it · the still-local-only **gap** · migration plan). Read this so a session pulls data from the bucket, not from someone's `C:/F:/Drive` ? the goal is to remove the local box as a blocker. Pairs with the two runbooks below.

31 -

../deploy/gcp/nightly_regression/README.md ? **? Nightly engine-regression OPERATIONS MANUAL (one-stop shop):** enable · run on demand · change any setting (email recipients/frequency/clips/config) · observe · pause/disable/remove. The full WASB-pipeline reel-count+cost+email nightly on an ephemeral GCP GPU VM (design/status: NIGHTLY_ENGINE_REGRESSION.md). The lighter CPU cached-trajectory tripwire is `scripts/nightly\_regression.py` (local Windows Scheduled Task, PR #202).

32 - TRACKNET_WSL_SETUP.md

33 - Evaluation harness and calibration drivers (in `backend/eval/`)

34

35 ## Data layer (collection · annotation · golden sets ? **no ML**)

36 - [data_pipeline/](data_pipeline/README.md) ? **the isolated data pillar:** how footage is recorded (`VIDEO\_RECORDING\_GUIDELINES`), golden sets are created + which videos (`GOLDEN\_SET\_IMPLEMENTATION` · `GOLDEN\_SET\_PHASE1\_VIDEOS` · `GOLDEN\_VIDEOS`), and data is shared/ingested (`GOLDEN\_DATA\_SHARING`). Deliberately **no segmentation/stitching/model** detail ? that's the ML pillar (top level + `COMPUTE\_DECOUPLED\_SERVING/`). The **annotation-vendor (Roora)** engagement is archived under `archives/roora-vendor/`.

37

38 ## Conventions & templates

39 - PROJECT_DOC_TEMPLATE.md ? Reusable skeleton for a **tracked project doc:** status header · decisions-locked · progress tracker (??/?) · definition-of-done · archived on completion. Standardizes the emergent pattern in `RALLY\_QUALITY\_RESEARCH.md` §7 and `GOLDEN\_REGRESSION\_FIXTURES.md`. Use for any substantial design/project work.

40

41 ## Historical documents

42 Completed plans, old research, ADRs, and checkpoints live in [`archives/`](archives/).

43

44 - ****June-2026 test-hardening sweep**** ? ****COMPLETED**** (2026-06-14; T1?T5 + full loop + verification per #162) and ****archived****; the former top-level pointer-stubs were removed 2026-07-02 (absorbed into the umbrella's ***Prior audits***). Full records: [`archives/CODE\_AUDIT\_AND\_TEST\_HARDENING-COMPLETED-2026-06-14.md`](archives/CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md), [`archives/DEFERRED\_HARDENING\_PLAN-COMPLETED-2026-06-14.md`](archives/DEFERRED_HARDENING_PLAN-COMPLETED-2026-06-14.md), [`archives/HARDENING\_LOOP\_HANDOFF-COMPLETED-2026-06-14.md`](archives/HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md).

45 - ****CODE_CLEANUP_AUDIT (2026-06-24)**** ? the prior deep engine audit (35 findings; Phase 1 complete, Phase 2 largely shipped, remaining = the `main.py` god-module split). ****Absorbed into the umbrella**** (see its ***Prior audits (absorbed)***) and archived ? [`archives/past\_projects/CODE\_CLEANUP\_AUDIT\_2026-06-24-SUPERSEDED-2026-07-02.md`](archives/past_projects/CODE_CLEANUP_AUDIT_2026-06-24-SUPERSEDED-2026-07-02.md).

46 - ****KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN**** ? the five-repo audit execution board, ****folded into the umbrella**** as its ***Remediation tracker (live)*** and archived ? [`archives/past\_projects/KHELSUTRA\_GURU\_HEALTH\_REMEDIATION\_PLAN-FOLDED-2026-07-02.md`](archives/past_projects/KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN-FOLDED-2026-07-02.md).

47 - DECISION_SNAPSHOT_REGRESSION.md ? ****?** Tracked project (decisions D1?D11 LOCKED 2026-06-25, #383; implementation P1?P7 pending): **** a deterministic **quality-neutral proof**** ? freeze + compare the pipeline's ****decision output**** (rally segments + stitch/concat plan) off a cached trajectory, GPU-free, in CI. The behavioral proof that lets a ***genuine*** (non-AST-identical) refactor be confidently called quality-neutral. ****Never hashes the encoded reel**** (non-deterministic). Extends the existing characterization snapshot (`test\_golden\_fixtures.py`) from segmentation to the stitch plan; peer of the M4 quality-floor + GPU `nightly\_reelcount.py`. S7 tracker is the source of truth.

48

49 - ****Field-ops / customer-feedback**** (recruiting flyer · field-visit playbook · PMF) ? [`archives/field-pmf/`](archives/field-pmf/) (archived 2026-06-21; PMF validation still an active ***business*** thread ? see `NEXT\_STEPS.md`).

50 - ****Roora annotation-vendor**** engagement (labeling brief · vendor bake-off) ? [`archives/roora-vendor/`](archives/roora-vendor/) (archived 2026-06-21).

51 - Superseded: per-video output layout ? [`archives/PER\_VIDEO\_OUTPUT\_LAYOUT-SUPERSEDED-2026-06-21.md`](archives/PER_VIDEO_OUTPUT_LAYOUT-SUPERSEDED-2026-06-21.md) (? `PER\_VIDEO\_WORKSPACE.md`).

52

53 See `archives/README.md` for the umbrella policy and `archives/past\_projects/README.md` for classification of small tech debt / starter projects.

54

24. user (2026-07-08T05:47:16.104Z)

13:[Omitted long matching line]

24:3b. ****[Claude] Five-repo audit health remediation ? ? COMPLETE (2026-07-02).**** All waves (P0?P28) landed and the last residuals were closed or filed as issues; the umbrella tracker was archived to `docs/archives/past\_projects/AUDIT\_REPORT\_khelsutra-guru\_2026-07-02-COMPLETED-2026-07-02.md`. Remaining owner-gated/future work lives as GitHub issues: #443 (CORS default), #444 (corpus Phase-2 upload), #445 (GCS eval plumbing), plus #301/#327/#332/#384.

153: rule, ground rules, Phase-2 record-and-show-back stretch goal) + `docs/FIELD\_VISIT\_ANSWER\_SHEET.md` (per-visit

25. user (2026-07-08T05:47:17.720Z)

12:<https://github.com/Khelsutra/badminton-highlight-indexer/issues/523>

16:1. Engine PR for #523: add an opt-in compile rendition field (`original` default,

24:4. Link each PR back to #523 and add verification notes/artifacts there.

26. user (2026-07-08T05:47:27.200Z)

1 # Next steps (live) ? refreshed 2026-06-21

2

3 > ****How to read:**** this top section is the ****current cross-track state + prioritized pending list**** (refreshed 2026-06-21, replacing the stale 2026-06-13 header). The detailed ****owned-model & rally-quality track**** plan (2026-06-13) is ****preserved below, still valid**** ? it is now ***one of several*** active tracks. Live blow-by-blow: memory ``current-state.md``. Doc lifecycle/hygiene: ``docs/DOC_STATUS.md``.

4

5 > ****? 2026-07-05 ? ACTIVE:** the control-plane durable rebuild.****** After the first owner browser CUJ

6 > blocked on the deployed CP (OOM/ephemeral-state ? DoD #1), the consolidated rebuild is specced,

7 > ****owner-confirmed****, and executing:
CONTROL_PLANE_DURABLE_REBUILD.md

8 > (\$7 there is the source of truth ? P1?P8 PRs, then the P9 rebuild+redploy+CUJ).

9

10 ## ? Where we are (2026-06-21)

11 ****DECOUPLED_COMPUTE** delivered on GCP + archived****** (#270); ****compute = GCP only**** (ADR-013, DO dropped). ****COMPUTE_DECOUPLED_SERVING**** subproject spun up (#272, ****ADR-014****) ? one pure core + local/cloud profiles; ****M0?M9 LARGELY SHIPPED**** (ADR-014 ratified 2026-06-21; M1?M4 #279?#283, M5 #286, M6 #287, M8 #288, M9-auth #290 all merged, default-OFF). ****Per-video workspace P0**** shipped (#264); ****doc-status register + archival due-diligence**** in place (#273). The ****khelsutra site + per-rally UI**** ship in parallel (sibling track). The ****rally-quality / owned-model**** track (below) is the core ML work, gated on ****growing the golden corpus****.

12

13 > ****? LAUNCH DECISION ? NO-GO** (owner, 2026-06-30).****** The rally-detection product is ****not**** cleared for an external/commercial launch; we stay in ****development****. This is the expected default for the dev phase. Two unblockers gate a future GO: **** (1) **** set + meet the ****ship-gate target (#172)**** ? currently unset, needs corpus growth (and the R2 tolF1 ablation, item 8); **** (2) **** resolve the ****WASB shuttle-detector weight provenance**** (commercial-clean licensing). P2b (#396) was a do-no-harm quality increment, ****not**** a launch gate ? merging dev PRs ? launching. Re-decide only when both unblockers clear.

14

15 ## ? Prioritized pending ? all tracks

16 ****P0 ? highest leverage / unblocks the most****

17 0. ****? [owner-priority, 2026-06-21] MAKE IT UI-DRIVEN + ALPHA-SHAREABLE ?** "a few recordings lying around ? highlights, from a page, no SSH".****** Today the cloud flow is ****SSH/CLI-only**** (``docs/CLOUD_GPU_DRYRUN_GUIDE.md``); ****M12 does NOT fix it.**** Needs (a separate khelsutra ``web/`` SPA + deploy track, NOT an engine seam): **** (i) **** deploy the engine as a reachable ****service**** (M7 ? Cloud Run service, ? owner GCP spend for a ***persistent*** endpoint), **** (ii) **** the ****SPA ingest/progress screen**** (khelsutra ``web/`` B-P2: enter/upload video ? ``/api/process`` ? poll ``/api/jobs`` ? reel; ZERO engine code), **** (iii) **** flip ****M9 edge-auth ON**** (merged #290, default-OFF) to gate alpha tenants, **** (iv) **** a ****compute PICKER** in the UI****** ("your machine vs our cloud", D13/D15). The engine seams make it ***possible***; the SPA makes it ***usable***. Full record: `[[ui-driven-cloud-serving-gap]]`. ****Concrete execution tracker after the n=15 corpus review:**** ``docs/ALPHA_LAUNCH_READINESS.md`` \$7.

18 0b. ****? [packaging track, IN PROGRESS ? 11 PRs merged 2026-06-21/22] Downloadable batteries-included app ?** ``docs/PACKAGING_PLAN.md``.****** The LIGHT-tier engine is bundle-ready: ``pip install`` ? ``indexer --app`` serves the ****bundled SPA**** single-origin, ****torch-free****, with app-home state / ffmpeg resolver / DNS-rebinding guard / single-instance lock / safe DB upgrades (P0a?e, P1a/b, P2-launcher/P2a/P2d/P2e). ****? Next (fresh focused pushes):**** **** (1) **** P3 ****GPU validation + screencast**** (L4, ~\$0.50, DELETE-after, \$100 cap ? ``[[gcp-vm-always-delete]]``; ? native-WASB 0-rally tuned-config blocker), **** (2) **** P2b-install (uv/script + truly-local default config), **** (3) **** P2c first-run wizard (cross-repo khelsutra SPA from ``/api/capabilities``), then P3-cloud (Cloud Run) + P4 (ONNX/train/annotate = north star). Memory ``[[packaging-app-plan]]`` / ``[[next-session-packaging-app]]``.

19 1. ? ****COMPUTE_DECOUPLED_SERVING` M0?M9 LARGELY SHIPPED**** (2026-06-21, ADR-014 ratified). M1?M4 (#279/#280/#282/#283), ****M5 #286, M6 #287, M8 #288, M9-auth #290**** all merged (default-OFF); the ****real M7 L4 run**** validated the worker (22 rallies, VM deleted, \$0). All \$8 gating Qs locked (D7?D17). ****Owner residual**** (not blocking Claude's default-OFF code): counsel ToS/DPA (M9-consent live + M11 live), real GCP prices+FX (M8 calibrate), Google OAuth app (M12 real), CF team-domain/AUD + ingress-lock; real cloud verification authorized within a ****\$100/?10k cap**** (D17).

20 2. ****[owner, ongoing] Grow the golden corpus**** ? multi-court badminton + ?3 matches/sport (TT first). The single highest-leverage move; the quality track ***and*** full-corpus training feed on it.

21 3. ****[owner + Claude . GPU] Full-corpus training (n?5)**** ? a real ``weights/<ver>`` (the

```

`0.1.0-l4` bundle is n=2 plumbing-only).
22
23     **P1 ? near-term, mostly Claude-doable (default-OFF, green-gated)**
24     3b. **[Claude] Five-repo audit health remediation ? ? COMPLETE (2026-07-02).** All waves
(P0?P28) landed and the last residuals were closed or filed as issues; the umbrella tracker was
archived to `docs/archives/past_projects/AUDIT_REPORT_khelsutra-
guru_2026-07-02-COMPLETED-2026-07-02.md`. Remaining owner-gated/future work lives as GitHub
issues: #443 (CORS default), #444 (corpus Phase-2 upload), #445 (GCS eval plumbing), plus
#301/#327/#332/#384.
25     4. **[Claude] `COMPUTE_DECOUPLED_SERVING` batch ? REMAINING: M11 + M12**
(M5/M6/M8/M9-auth ? merged this session). **M11** = `backend/flywheel.py` capture?store?shadow-
eval?goldenize plumbing reusing `workspace.py::for_video` + `backend/eval/experiment.compare` +
`backend/eval/promotion.py::promote_if_served_safe`, behind `flywheel.*`=OFF + a faked
`consent_attested` (? privacy-sensitive ? give it fresh focus). **M12** =
`backend/storage/gdrive_api.py` `GDriveApiStorage(StorageBackend)` + a `gdrive-api://` scheme
(regex already allows the hyphen; add to `_KNOWN` + `_backend_class` + the StorageConfig
Literal) against an INJECTED fake Drive (mirror `gcs.py` DI); keep the mount-only `gdrive://`
untouched. ? **M9-consent cols = a new v7 migration** (M8 took v6). M9-consent ToS/DPA text =
owner/counsel.
26     5. **[Claude] Per-video workspace P1?P6** (P1 v6 DB migration ? P2 worker ? P3 per-video
routing of `detector/+`tmp/` * [the `output/chunks_` de-hardcode base = ? done via #282]* ? P4
? P5 Launch-Report flip ? P6 GPU parity). ? P1's v6 migration coordinates with the M8 v6 cost-
ledger migration ? author the version bump once.
27     6. ? **Ops-Agent codified (#291)** ? `deploy/gcp/create_gpu_vm.sh` (the one-command VM
helper) always bakes the Ops-Agent startup-script; provision.sh grants the APIs+SA roles. PROVEN
active on a fresh VM 2026-06-21. ([[gcp-ops-agent-enable]])
28     7. **[owner] Set the ship-gate target (#172)** ? needs corpus growth.
29     8. **[owner-side GPU] R-series: promote-or-reject `R2` toLF1 as the gate** (the
ablation) ? needs GPU-extracted golden features.
30
31     **P2 ? later / gated / deferred**
32     9. Rally-quality **P3/P4** (person-fusion LOV0, audio cue) ? needs GPU golden features.
33     9b. ? **DONE (#396):** the R4 density fix (finding B / P2b in
`RALLY_DETECTION_FIXES_PLAN.md`) is **promoted to default-ON**. Golden-set A/B at the **SERVED**
windowing showed toLF1 0.353?0.364 (?toLF1 +0.010, no regression), so
`inactivity_temporal_density` default was flipped to True. (Landed flag-gated default-OFF in
#394.) Fixed the sparse-tracking quirk where one fast sample after a detection gap held the
rally tail open.
34     10. **M0.4 scorer swap** ? TemporalMaxer (per the 2026-06-17 external review).
35     11. **Multisport** (table-tennis first, then pickleball).
36     12. Serving **M8/M9** ? per-video cost ledger, edge auth, DPA/consent (within
`COMPUTE_DECOUPLED_SERVING`; owner/legal-gated).
37     13. **WASB-C3** weight provenance (gates *sharing* a trained model) ? legal.
38     14. **Rename repo off "badminton"** (owner, BACKLOG).
39     15. Owner housekeeping: retire the shared **DO droplet** + **DO token**; collector
**`md5`** export (collector repo).
40     16. **PMF field validation** (owner-side, business track).

```

27. user (2026-07-08T05:48:09.912Z)

```

1     # DEPLOY_REPORT ? cloud-serving (M7) first real L4 run + GPU optimization sweep
2
3     > **Date:** 2026-07-04 . **Owner-authorized** GCP spend (D17, $300 envelope) . deleted
every
4     > resource after (`[[gcp-vm-always-delete]]`; $5 residual audit = $0 of ours).
5     > **Video:** `MahadevpuraSingles.MP4` (owned golden, 13,418 frames). **Job:** Cloud Run
L4
6     > (Ada sm_89, 24GB) . 8 vCPU / 32GB . `--no-gpu-zonal-redundancy` . scale-to-zero.
7     > **Mode:** fully-local (`skip_ai_handoff=true`) ? the number is the **owned WASB path
only, zero Gemini**.
8
9     ## 1. M7 infra proven
10    - Worker image builds + pushes (`asia-southeast1/rally`); the M6 `CloudRunCompute` shim
dispatches the
11    L4 Job and bucket-polls the done-marker; WASB weights via a GCS-FUSE read-only mount

```

```

(no weights in
12     image ? WASB-C3); `intervals.csv` + `report.json` land in the bucket.
13     - Owned model runs on the L4: `native exec wasb_infer.py`, 13,418 trajectory points, 48
windows, 0 Gemini.
14
15     ## 2. GPU optimization ? measured
16     **Headline: the L4 is CPU-FEED-bound, not GPU-bound.** GPU utilization sat **<1%** on
every config ?
17     a GPU that idle is *waiting*, not computing. So kernel/GPU levers do nothing; the wins
(and the
18     remaining headroom) are all on the CPU feed.
19
20     | # | Config | Phase-2 wall | vs baseline | GPU util | Parity (rally windows) |
21     |---|-----|-----:|-----:|:-----:|-----:|
22     | 0 | cu113, batch 8 (baseline) | 727.9s | ? | <1% | reference |
23     | 1 | torch 2.2.2+cu118, batch 8 | 722.1s | ?0.8% (noise) | <1% | intervals byte-
identical; trajectory differs (6780 vs 6777 pts) |
24     | 2 | cu113, **batch 64 + prefetch 4** | **506.1s** | **?30.5% (1.44x)** | ~1% | **byte-
identical** ? (6777 pts) |
25     | 3 | cu113, batch 128 + prefetch 6 | OOM | ? | ? | exceeds 24GB VRAM (`CUDA OOM, 9
GiB`) |
26
27     ### Dispositions
28     - **? SHIP ? batch/prefetch tuning (config #2):** 30.5% parity-safe. Batch 64 is at the
L4 VRAM
29     ceiling (128 OOMs), so it's the max viable value, not arbitrary. ? PR `claude/wasb-
gpu-batching`.
30     - **? DISCARD ? torch cu113?cu118 arch bump (config #1):** 0.8% = noise (JIT was never
the wall,
31     since the GPU was idle regardless), *and* it perturbs trajectory numerics (a byte-
parity risk).
32     Empirically dead for this workload.
33     - **? DON'T ? bigger GPU:** idle at <1% ? an A100/H100 would also idle. Pure waste (D8).
34     - **? OPT-IN ONLY (never default) ? FP16 / TF32 / cudnn.benchmark / NVDEC:** each
changes detector
35     numerics ? breaks the D5/D16 byte-parity guarantee (verified reasoning in the lever
workflow).
36
37     ## 3. Root cause & the remaining ~5x (scoped follow-up)
38     Streaming forces `DataLoader(num_workers=0)` (the in-memory `SequentialFrameStore` +
`read_image`
39     shim can't pickle to worker processes), so decode?PIL?ToTensor runs on **one core**,
inline with the
40     GPU; sliding windows re-transform each frame ~`frames_in`x. At batch 64 the GPU is
*still* ~1% ?
41     batching only removed the per-batch round-trip/Python overhead; the ~500s floor is the
single-threaded
42     feed. **The remaining 3?5x is a parallel-feed change**, not a config or a bigger
machine:
43     - **Design:** partition the video into contiguous chunks, one DataLoader worker per
chunk (each opens
44     its own `VideoCapture`, seeks to its chunk start, reads forward ? sequential within a
chunk,
45     parallel across chunks); collate reorders. Parity-safe (same frames, same transform,
same order).
46     - **Alternative (rejected for Cloud Run):** the existing disk-frames path already does
`num_workers>0`,
47     but materializing 13,418 PNGs into the RAM-backed `/tmp` (~27GB) OOMs the 32GB box ?
which is *why*
48     streaming exists. A real scratch disk (GCE VM, not Cloud Run) would unlock it.
49     - Only after the feed is parallel does a **higher-vCPU config** pay off (more cores to
feed the GPU).
50
51     ## 4. Cost ledger (D17)
52     - Builds: 4 x ~$0.16 (cu113 base reuse, cu118, batch) ? $0.5

```

```

53 - Runs: baseline + cu118 + batch64 + batch128(OOM, short) + probes ? ~$2?3 of L4 time
54 - Idle: **$0** (scale-to-zero; GPU bills only during execution)
55 - **Total well under the $300 / D17 envelope.**

```

28. user (2026-07-08T05:48:10.556Z)

```

backend\orchestration.py:53:_LOCAL_COMPUTE_LANE = threading.BoundedSemaphore(1)
backend\orchestration.py:918:          # Serialize the CPU/GPU-heavy in-process detection (see
_LOCAL_COMPUTE_LANE). A remote
backend\orchestration.py:920:          with _LOCAL_COMPUTE_LANE:
backend\orchestration.py:1163:          # dispatch) doesn't run multiple compiles at once on the
box (see _LOCAL_COMPUTE_LANE).
backend\orchestration.py:1164:          with _LOCAL_COMPUTE_LANE:
backend\api\job_worker.py:37:# orchestration._LOCAL_COMPUTE_LANE regardless of the worker count,
so the single-box invariant holds
backend\api\job_worker.py:76:    stitch serialize via orchestration._LOCAL_COMPUTE_LANE
regardless of the worker count ? so the

```

29. user (2026-07-08T05:48:18.438Z)

```

1      # Testing Strategy ? Compute-Decoupled Rally Serving
2
3      > **Companion to** [`README.md`](README.md) ($7 tracker is the source of truth).
**Status:** DRAFT (2026-06-21).
4      > **Thesis:** the decoupling is only *real* if we can **prove the core is unaffected
across profiles** ? for **$0**, with no GPU/WSL/network** ? and quarantine the one thing that
genuinely can't be byte-deterministic (cross-GPU detection) to an owner-side real-hardware
layer. Tags: `[verified]` = building block already in the tree; `[design]` = to build.
5
6      ---
7
8      ## 0. The two guarantees we are testing
9      1. **Decoupling guarantee** ? the same `(input bytes + config)` produces the **same
windows + the same stitched segment ranges** whether the core runs truly-local or on Cloud Run.
*Provable for $0* because detection can be replaced by a deterministic `replay_csv` stub.
10     2. **Quality guarantee** ? real WASB detection on a real GPU is *good enough at the
rally level*. **Not** byte-deterministic across GPU architectures (L4 vs 2070 differ sub-pixel ?
684/1195 frames, per the 2026-06-20 deploy report), so it is judged at the **window/rally
level**, owner-side, on the golden corpus.
11
12     The whole strategy hangs on keeping these **separate**: CI gates #1 deterministically;
the owner gates #2 on real hardware.
13
14     ## 1. Test layers
15
16     ### L1 ? Unit (pure core) `[verified]`
17     - **What:** `trajectory_to_action_windows` determinism; `_merge_overlapping`;
`_watermark_filter` string build; `_parse_time`; **profile preset/precedence resolution** (new);
`StorageRef.parse_uri` scheme dispatch.
18     - **How:** pytest, no GPU/net/WSL; same-input?same-output; `abs-tol 1e-6` on floats.
19     - **Modes:** all (pure fns are profile-independent).
20
21     ### L2 ? Contract (seam ABCs) `[verified]`
22     - **What:** `DetectorRunner` ABC ? `run_predict` CSV schema + `healthcheck` non-raising
? across `Wasb`/`NativeWasb`/`Stub`; `StorageBackend` ABC roundtrip (`put` ? exists ? stat ? list
? fetch ? delete) across local/s3/gcs/gdrive; the new `DeviceContext` probe contract.
23     - **How:** stub runner for detection; **injected fake S3/GCS clients + gdrive=`tmp_path`
mount** (`tests/test_storage.py`); assert the CSV columns are identical across runners.
24     - **Modes:** `cpu-mock` proves the contract; local/cloud use the same code paths.
25
26     ### L3 ? Integration (GPU-free E2E) `[verified]`
27     - **What:** the full trajectory-hybrid pipeline (detect ? window ? persist ? stitch) on
the stub detector; `worker cmd_infer` fetch ? segment ? put with fake storage; chunked-upload
assemble.
28     - **How:** `RALLY_DETECTOR_IMPL=stub` (synth **and** `replay_csv`); storage fakes;

```

```

`tmp_path` workspace; assert `intervals.csv` + `report.json` + a reel are produced
deterministically.
29     - Modes: `cpu-mock`, but it exercises the cloud-serving code path
(fetch/put/dispatch) with no cloud.
30
31     L4 ? Profile-PARITY (the load-bearing proof) `[design]`
32     - What: the explicit decoupling guarantee ? run the same `replay_csv` golden
trajectory through (a) the in-process truly-local path and (b) the worker + fake-`gcs`
(cloud-serving-mock) path, and assert byte-identical candidate windows + identical stitched
segment ranges (and, on a fixed clip, an identical reel via ffmpeg determinism).
33     - How: one test parametrized over `{profile=truly-local, profile=cloud-serving-
mock}`, fed the same replay CSV + same windowing config; diff the window list and segment
ranges exactly. Because detection is a deterministic replay here, any divergence is
necessarily a profile/seam leak ? exactly what must never happen.
34     - Modes: truly-local vs cloud-serving (mocked). This is the test that, if green,
proves the whole design.
35
36     L5 ? E2E real-GPU (owner-side) `[design]`
37     - What: real WASB detection on a real GPU (local 2070 ? truly-local; GCP L4 ? cloud-
serving) end-to-end ? reel; cross-GPU parity checked at the rally/window level (NOT CSV
bytes).
38     - How: owner runs the golden corpus on the GPU box and on Cloud Run; compare
rally counts/windows; smoke-test the reel (`ffmpeg -f null - full decode exit 0 + a 5-frame
contact sheet visually confirming real rallies). Recorded in a DEPLOY_REPORT under
`runlogs/`.
39     - Modes: truly-local + cloud-serving on real hardware.
40
41     ## 2. Real-GPU vs mocked ? the split
42     - Mocked is the CI default. The stub detector (CPU, deterministic, byte-exact via
`replay_csv`) + in-memory storage fakes prove all plumbing + the profile-parity guarantee
with $0 and no GPU/WSL/network.
43     - Real-GPU is owner-side only (the golden corpus + labelled videos live with the
owner). It validates that NativeWasbRunner actually uses the GPU and that cross-GPU detection
diffs are sub-pixel ? judged at the window level, never CSV bytes.
44     - The contract: CI gates the deterministic post-trajectory transforms (window +
stitch) byte-exact; real-GPU gates that detection on the chosen GPU is good enough at the
rally level. Honest-limits caveat (golden $6) baked in: green CI proves
plumbing/determinism, not field quality.
45
46     ## 3. CI gates (extend the existing lanes)
47     Today's lanes: `ruff`, `mypy` (lenient baseline), import-smoke (no GPU stack), full
suite (offline, mocked, `--cov-fail-under`), golden-fixtures cross-OS (ubuntu/windows/macOS).
Add:
48     - Profile-parity lane ? the L4 test (truly-local vs cloud-mock identical
windows/segments).
49     - Storage-fakes roundtrip lane ? L2 across all four backends.
50     - Per-PR core-touching gate ? any PR touching detect/window/stitch callers (esp.
M3, the `output/` de-hardcode) must show: golden fixtures green on all 3 OSes AND
`experiment.compare(CONTROL, profile-variant)` F1-delta ? 0 (no-regression), AND coverage
not regressed vs base.
51     - No real GPU / net / WSL in CI ? ever. Real validation is the owner's L5.
52
53     ## 4. Profile-parity ? why it's airtight
54     The decoupling claim is proven by feeding one `replay_csv` golden trajectory through
both the in-process path and the worker+fake-bucket path and asserting byte-identical
windows + segment ranges. Detection is replaced by a deterministic replay, so the only
variable is the seam (compute target + storage backend). If the outputs ever diverge, a
profile leaked into the core ? the precise failure D1 forbids. Real cross-GPU detection diffs
are quarantined to L5 and judged at the window level.
55
56     ## 5. Sample runs / runlogs for posterity (the DECOUPLED convention)
57     Each real run writes a dated report under [runlogs/](runlogs/) (e.g.
`DEPLOY_REPORT_{profile}_{date}.md`), mirroring
`DECOUPLED_COMPUTE/DEPLOY_REPORT_GCP_2026-06-20.md`, capturing:
58     - resources provisioned + the exact commands;

```

```
59      - run results ? **rally counts, parity numbers**, `gpu_active_seconds` / `wall_seconds`  
/ `bytes_out`;  
60      - a **reel smoke-test** ? `ffmpeg -f null -` exit 0 + a 5-frame contact sheet visually  
confirming real rallies;  
61      - gotchas (conda-ToS, WASB-C3, GPU stockouts, font fallback) + the **final teardown  
state** ($0 audit);  
62      - so the next deploy is **one-shot**.  
63  
64      **DoD:** `truly-local` (M5) and `cloud-serving` (M7) each get **?1 committed sample  
run** here.  
65  
66      ## 6. Coverage map ? what each milestone must prove before merge  
67      | Milestone | Must be green |  
68      |---|---|  
69      | M1 config | L1 preset/precedence; profile='' = today (no behaviour change) |  
70      | M2 input_backend | L2 storage roundtrip (4 backends); L3 worker fetch with  
local=identity |
```

30. assistant (2026-07-08T05:49:37.457Z)

I have everything needed. Compiling the structured code map for the DOCS + CONVENTIONS subsystem.

31. user (2026-07-08T05:51:52.379Z)

Structured output provided successfully