

claude-windows / e663947f-2127-4383-9905-650340ff9863

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\`e663947f-2127-4383-9905-650340ff9863.jsonl`  
- SHA-256: `5348cea02bb1f124cf679bbe5b7a664faadf3948f102aadd2a231987bbe38f5`  
- Source modified: `2026-07-09T08:10:15+00:00`  
- Imported at: `2026-07-09T08:30:30+00:00`  
- project: `C--Users-avidu-Projects-badminton-highlight-indexer`  
- session_id: `e663947f-2127-4383-9905-650340ff9863`
```

Transcript

1. user (2026-07-09T06:23:48.042Z)

The cloud-serving worker infer path (`backend/worker/\_\_main\_\_.py`) has three linked observability gaps, surfaced during a live production CUJ on 2026-07-09 (video GX010155, video_id `e2e144175c8d05451903ae44a9ce0aa0`, worker execution `rally-worker-dpssz`, deployed image `sha256:1838c07f?` from `origin/master @ 00e8f92`). Fix all three in ONE PR (branch from `origin/master`).

CONTEXT: `phase\_placement.validation/segment/compile` are all `cloud\_run\_l4`, so the control-plane delegates validation to the L4 worker and dispatches it WITHOUT `--skip-validation` (verified in the execution args). The worker therefore DOES run the full validator suite and re-computes the checksum ? but none of it is observable.

GAP 1 ? validation invisible on a PASS. In `backend/worker/\_\_main\_\_.py` (~lines 281?343) the worker runs `validator\_registry.run\_all\_with\_context(...)` unless `--skip-validation`, but only `print()`s on `[worker] validation SKIPPED` or `[worker] validation FAILED`; a PASS logs nothing. And results are persisted only to the ephemeral worker DB via `db.add\_video(video\_id, ..., [r.model\_dump() for r in val\_results])` ? they are NOT in the pushed `report.json` (which contained only `video\_id, status, segment\_count, video\_uri, segments`). So a passing validation is invisible in both logs and the runlog. Users legitimately thought validation was skipped.

GAP 2 ? no per-stage wall-clock timing. There is no timing around checksum (`compute\_video\_id`), validation (`run\_all\_with\_context`), or segmentation (`segmenter.process\_video`). In the CUJ, segmentation took ~10 min with nothing recorded.

GAP 3 ? the #534 live progress heartbeat + telemetry does NOT fire on the native-L4 path. Evidence: over a ~10-min GPU pass, `gs://kheutra-rally-serving/outputs/<video\_id>/run\_telemetry.jsonl` contained ONLY its meta header (116 bytes: `{ "meta": { "label": "l4-infer-?", "gpu\_name": "NVIDIA L4", "cpu\_count": 8, ... } }`) with ZERO snapshots; there were NO `[worker] progress:` lines in stdout; and no `<video\_id>\_detector\_progress.json` was ever written to GCS. The snapshots + heartbeat are driven by the detector progress callback `\_progress` (`backend/worker/\_\_main\_\_.py` ~line 388, passed to the segmenter as `progress=\_progress` ~line 434), which was never invoked. Root-cause in `backend/pipeline/detectors/native\_wasb\_runner.py` (~lines 284?440): it is meant to stream `wasb\_infer` stdout and parse the periodic `done/total (pct%)` lines (emitted every `log\_every\_batches`=50) to call the callback ? verify/repair that path so the callback fires ~every `detector\_progress\_sample\_sec` (15s).

FIX (single PR):

1. Clear per-stage log lines in the worker infer flow, emitted even on success: `[worker] checksum: video\_id=? (Xs)`; `[worker] validating (N checks)?` then `[worker] validation PASSED/FAILED in Ys ? <name>=ok/FAIL(score)?`; `[worker] segmentation (<segmenter>) done in Zs ? R rallies`.
2. Time each stage with `time.perf\_counter()` and persist BOTH a `timings` block AND the validation results (per-validator name/passed/message/score) into the pushed `report.json` (built/pushed via `\_serialize\_and\_push\_outputs`, ~lines 125?158) so runlogs capture time-per-stage + validator outcomes.

3. Make the native-L4 detector progress callback actually fire so `run\_telemetry.jsonl` gets real GPU/velocity snapshots, the `[worker] progress:` heartbeat appears in stdout, and `

VERIFY: add offline unit tests (load the module by file path, mirroring `tests/test\_nightly\_regression.py`) for the timing/report-serialization additions and the progress-line parsing. Run `python -m ruff check .` (install ruff==0.15.16, CI's pinned version), `mypy backend training`, and `python run\_all\_tests.py --cov=backend --cov=training --cov-fail-under=70`. NB: on Windows there are 2 pre-existing `tests/test\_storage.py` GCSFuse failures unrelated to this work (fixed separately in #548) ? they pass on Linux/CI. The live-L4 heartbeat itself needs an owner-driven CUJ to fully confirm end-to-end; state clearly what was and wasn't verified. Open a PR to master.

2. assistant (2026-07-09T06:23:55.245Z)

I'll start by getting oriented ? the memory tells me the primary working directory is a stub and the real code is elsewhere. Let me confirm the actual repo location and read the session handoff.

3. user (2026-07-09T06:23:56.311Z)

<system-reminder>This memory is 2 days old. Memories are point-in-time observations, not live state ? claims about code behavior or file:line citations may be outdated. Verify against current code before asserting as fact.</system-reminder>

```
1      ---
2      name: actual-repo-location
3      description: "The badminton-highlight-indexer code lives under khelsutra-guru/, not the
primary working dir stub"
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: df5d5fb2-534a-447a-8181-b222f94655f3
8      ---
9
10     The primary working directory `C:\Users\avidu\Projects\badminton-highlight-indexer` is a
near-empty
11     **stub** ? it contains only `AGENTS.md`, an empty `.claude/`, and
`output/nightly_regression/`
12     (a `LATEST_STATUS.txt`). The actual `khelsutra-guru/badminton-highlight-indexer`
codebase lives at:
13
14     `C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer`
15
16     (git remote `github.com/Khelsutra/badminton-highlight-indexer`, default branch
`master`).
17
18     **How to apply:** when a task references this repo's code (backend/worker,
pipeline/segmenters,
19     config.json, run_all_tests.py, etc.), `cd` to the khelsutra-guru path and work there ?
the stub dir
20     has none of it. Confirmed 2026-07-06 while fixing the wf294 silent-detector-failure
issue.
21
22     Repo conventions verified there: CI = ruff 0.15.16 (`ruff check .`), mypy 2.1.0 (`mypy
backend
23     training`; `mypy.ini` quarantines several modules incl. `backend.worker.__main__`,
`gemini`,
24     `yolo_hybrid` via `ignore_errors=True`), full suite via `python run_all_tests.py`,
coverage floor
25     `--cov-fail-under=70` (actual ~86%). Segmenters use a `Success`/`Failure` Result
contract
26     ([[efficient-timeouts-and-retries]] is unrelated).
27
```

4. user (2026-07-09T06:23:57.327Z)

```
1      ---
2      name: session-handoff-gpu-resident
3      description: "Start here to pick up the GPU-resident WASB rollout (issue #506) ?
implementation PR #507 is open; golden gate next"
4      metadata:
5          type: project
6      ---
7
8      # Session Handoff ? GPU-resident WASB inference (rollout) . issue #506 . PR #507
MERGED
9
10     ## LATEST 2026-07-07 (end of session) ? all serving fixes merged + fresh redeploy
11     Every gap surfaced this session is merged to `origin/master @ ef41473`: **#511**
(compute=local fails fast on the weightless control-plane) + **#516** (SPA compute-picker
defaults to CLOUD on hosted) = the compute-picker footgun; **#512** (validators share ONE decode
pass ? measured ~3.1x on the 5.3K clip: warm 179s?58s on a 16-core box; projects ~6.5min?~2min
on the 2-vCPU control-plane); **#513** (worker skips redundant re-validation on dispatch ? also
fixes the default-blur-threshold reject); **#515** (worker no-done-marker fastfail). **Clean
rebuild from `ef41473` ? redeploy control-plane + worker to the new digest (D3)** ? see the
DEPLOY_REPORT / [[gpu-vm-setup-gotchas]]. #509 forwarding was PROVEN live by the CUJ (worker
args carried `decode_backend=nvdec_resident`+batch64+prefetch4). Two follow-ups SPAWNED:
**task_5acf4b46** (clean 2?3 video validation-speedup closure test + DEPLOY_REPORT) and
**task_4585aa00** (cache validation results by video_id ? skip re-validation on reupload; the
missing cache that made CUJ retries re-pay validation). Validation is NOT cached across
reuploads today (only a whole-pipeline cache gated on a prior SUCCESSFUL segmentation).
12
13     ## SHIPPED 2026-07-07 ? GPU-resident NVDEC decode is LIVE in cloud-serving
14     Full arc done: PR #507 (feature) + PR #508 (cloud-serving preset default =
nvdec_resident) both merged to master (44a2f47, 143dbb9). The `rally-worker` Cloud Run Job
(asia-southeast1) runs the new image `rally-worker:latest` `sha256:465a9250?` (torch 2.6 +
PyNvVideoCodec 2.1.0 + baked NVDEC libs) with `WASB_DECODE_BACKEND=nvdec_resident` set ? **nvdec
is the live serving default**, validated on the real L4 (GX010128 served trajectory reproduced
the VM: 26479?26481 visible; cpu+nvdec+canary smokes all ?). Record:
`docs/COMPUTE_DECOUPLED_SERVING/runlogs/DEPLOY_REPORT_cloud-serving_gpu-resident-
nvdec_2026-07-07.md`.
15     - **ROLLBACK levers:** `gcloud run jobs update rally-worker --region asia-southeast1
--remove-env-vars WASB_DECODE_BACKEND` (? torch-2.6 cpu path, same image); or `--image ?/rally-
worker@sha256:3bfefaf6?` (? prior torch-1.11 image).
16     - **PRs:** #507 (feature) + #508 (preset default) + #509 (dispatch forwarding + release
hardening) all merged. master @ 66756ad.
17     - **DONE ? control-plane + worker both on image D2 `sha256:09d2edf7` (built from master
66756ad = #507+#508+#509).** Redeployed 2026-07-07: `gcloud run services update rally-control-
plane --image ?:latest` ? rev **00017-qbd** (100% traffic, booted healthy: uvicorn up, cloud-
serving profile ACTIVE, edge-auth on); `gcloud run jobs update rally-worker --image ?:latest`.
So the #509 dispatch **forwarding** + #508 preset default are now LIVE in the control-plane
image. NB the FIRST control-plane deploy this session used the stale `465a9250` (built from
44a2f47, pre-#508/#509) ? rev 00016; superseded by 00017. Rollback: `gcloud run services update-
traffic rally-control-plane --to-revisions rally-control-plane-00015-kp9=100` (last pre-nvdec
rev) or `--to-revisions ?-00016-?`.
18     - **`WASB_DECODE_BACKEND=nvdec_resident` env is still set on the worker Job** ? now
REDUNDANT with the forwarded --set + preset (both say nvdec, env just wins), kept as belt-and-
suspenders + manual override. Safe to remove ONLY after a real /api/process dispatch confirms
the forwarding path live (needs a CF Access JWT ? not verifiable this session; forwarding is
deployed + unit-tested but not live-dispatch-verified).
19     - `DEPLOY_REPORT?gpu-resident-nvdec_2026-07-07.md` is current: PR #510 added the $5
"Update" (D2 redeploy of control-plane rev 00017 + worker) and corrected $6. All four PRs
merged: #507 #508 #509 #510; master @ 3be0935.
20     - **Wrinkles fixed (#509):** dispatch now forwards
indexing.wasb.{decode_backend,batch_size,prefetch_batches} (worker has no DEPLOYMENT_PROFILE ?
never applied the preset; batch tuning was silently lost too); added `.gcloudignore` (Cloud
Build had none ? relied on the .gitignore fallback to skip 8 GB output/scratch); added
`deploy/cloudrun/post_deploy_smoke.ps1`. See [[gpu-vm-setup-gotchas]].
21     - WASB-C3 weights-license gate still open (unchanged; pre-existing).
```

```

22
23     ## Live CUJ 2026-07-07 ? release PROVEN, 3 pre-existing gaps surfaced (all filed as
spawned tasks)
24     Owner ran a real /api/process CUJ ("Video 1 - Badminton.MP4", a 5312x2988 **10-bit
HEVC** 120Mbps clip). Outcomes:
25     - **#509 forwarding PROVEN live**: the cloud dispatch put `--set
indexing.wasb.decode_backend=nvdec_resident --set ?batch_size=64 --set ?prefetch_batches=4` into
the worker exec args. And a sharp-clip green run on the D2 worker (exec `rally-worker-ql56s`,
env-driven nvdec) = 19 rallies, success. Serving path is GREEN for valid footage.
26     - **Gap 1 ? slow validation (task_dcee8bf0)** ~125 random `cap.set(POS_FRAMES)` seeks
across scene_cut(100)/blur(15)/relevance(10), each their own VideoCapture, on 5.3K 10-bit HEVC
on the CPU-only control-plane ? ~6.5 min. Fix: single sequential/keyframe pass.
27     - **Gap 2 ? compute=local footgun (task_81572190)** SPA sent compute=local ? in-process
on the weightless control-plane ? "weights not found" fail. Guard/redirect it.
28     - **Gap 3 ? worker re-validates with DEFAULT thresholds (task_72d98f83)** control-plane
min_blur=8.0 (baked) NOT forwarded; worker re-validates at default 20.0 and REJECTED this soft
clip (sharpness 11.9) ? rc2. This is what failed the CUJ on the cloud path. Fix: skip worker re-
validation on dispatch (control-plane is authoritative) or forward validation config.
29     - **Validation-latency design (my rec)** don't move validation to GPU standalone. (1)
stop the double-validation, (2) kill the seeks (sequential/keyframe), (3) end-state = ffprobe
metadata pre-gate on control-plane + fold content checks (blur/court) INTO the worker's existing
GPU decode pass (decode once, validate+detect together). All 3 gaps are pre-existing, NOT caused
by the nvdec release.
30
31     ## (historical) release/rollout phase ? superseded by SHIPPED above
32     PR #507 squash-merged to master as **44a2f47** (branch deleted). Pre-merge gate
reproduced locally: full suite 1764 passed / coverage 85.02% / ruff+mypy clean; CI 11/11 green;
2 optional review follow-ups landed (PyNvVideoCodec==2.1.0 pin in both deploy files;
decode_backend persisted in native-runner detection_params). Real-L4 validation + golden gate
PASS + latency/resource A/B all done (see PR #507 comment + below). **Remaining = the serving
release (image rebuild + gated nvdec flip); decode_backend still defaults cpu so nothing in prod
has changed yet.**
33
34     ### Serving release plan (live infra: Job `rally-worker` + service `rally-control-
plane`, asia-southeast1; ONE image `rally-worker:latest` drives both ? gen_service_yaml.py:193)
35     - **Phase A (ship image, no behavior change)** `gcloud builds submit
--config=deploy/cloudrun/cloudbuild.yaml --project gen-lang-client-0306026826 .` ? new `rally-
worker:latest` (torch2.6+pyncv2.1.0+NVDEC libs). Push does NOT change prod (Job/service pin a
digest). Cut over via `gcloud run jobs update rally-worker --region asia-southeast1 --image
?:latest` + redeploy service (`gen_service_yaml.py` ? `services replace`). decode_backend
defaults cpu ? cpu path under torch2.6 (validated F1 0.576?0.582). Smoke 1 clip. Rollback: `jobs
update --image <old digest>`.
36     - **Phase B (enable nvdec, gated + reversible)** flip via WORKER env
`WASB_DECODE_BACKEND=nvdec_resident` (native runner reads env first ? no code/preset change,
instant rollback by removing it). ? THE untested risk: NVDEC on the real Cloud Run L4 ? the
image bakes libnvcuvid 580.159.03; Cloud Run's host-driver ABI is unverified there (only tested
on a self-managed L4 VM). MUST smoke 1 clip on the Job before trusting it. Then canary; then
durable default = add `decode_backend: nvdec_resident` to the cloud-serving preset (presets.py)
+ rebuild + Launch Report.
37     - Cost: build ~$0.3; each Job smoke ~$0.4?0.7 (L4). Rollback levers at every step.
38
39
40     ## You are here
41     The implementation is DONE and **PR #507 is open** (branch `feat/506-gpu-resident-
nvdec`, 2026-07-07): default-OFF `indexing.wasb.decode_backend = "cpu" | "nvdec_resident"`, the
validated NVDEC?affine-`grid_sample`?GPU-normalize path wired into
`backend/pipeline/detectors/wasb_infer.py` (`NvdecResidentFrontend` +
`run_video_nvdec_resident`, same resumable cache keyed by decode_backend), native-runner
threading + healthcheck gates (CUDA + PyNvVideoCodec), and the env bump (wasb env ? py3.10 +
torch 2.6.0+cu124 + PyNvVideoCodec; libnvcuvid/libnvidia-encode staged in
`deploy/cloudrun/Dockerfile` + `deploy/digitalocean/gpu_setup.sh`). Verified locally: 1764 tests
green, coverage 85% (floor 70), ruff/mypy/link/leak clean; a 14-agent adversarial review
confirmed 7 minor findings, all fixed (notably `padding_mode="zeros"` = cv2 BORDER_CONSTANT for
non-16:9 sources). Research evidence unchanged: [[gpu-resident-modernization]], issue #506
comments (F1 0.574 ? baseline 0.582 on GX010128 at SERVED, ~1.7x, GPU 47?79%).

```

```

42
43     ## Real-L4 validation ? DONE + golden gate PASS 2026-07-07 (this session, ~$3 of the
$100 budget; VM torn down)
44     Validated end-to-end on an L4 (`rally-506`, torn down clean, no orphan disks) ? full
evidence in the PR #507 comment. All GREEN: (1) `gpu_setup.sh` builds the env clean (py3.10 +
torch 2.6.0+cu124 cuda True + PyNvVideoCodec 2.1.0 + NVDEC lib extraction); (2) on-box **parity
gate PASS** (`pytest -k parity`, 480s ? needed test-name fix 9267255, the documented `-k parity`
selected 0 tests); (3) crash/resume byte-identical; (4) cpu?nvdec cache isolation correct; (5)
**Cloud Run image builds + `--gpus all` smoke PASS** (baked NVDEC libs decode; native
healthcheck ok with nvdec_resident); throughput 43.7 fps / GPU 82% on GX010128.
45     **Golden gate @ SERVED = PASS** (10 baseline-comparable clips, 445 rallies): mean
?f1@0.5 **+0.020**, mean ?tolF1 +0.006, worst ?0.009 (**GX010128 must-pass: 0.574 vs 0.582 ?**),
8/10 improved-or-flat. The other 5 golden clips are 4K-sourced (baselines are 1080p) so a vs-
baseline number would conflate resolution with decode ? the parity gate already covers nvdec?cpu
on identical pixels, so they add nothing to the verdict.
46     **Latency + resource A/B** (2nd L4 `rally-506b`, torn down; GX010128, batch 8, telemetry
in `gs://?nightly/dev/506/results/bench/`): nvdec **43.7 fps @ GPU 83% / CPU 30%** vs cpu-
stream(no-prefetch) 15.0 fps @ GPU 28% / CPU 49%; GPU mem +170 MiB (of 24GB), RAM ~2?3 GB flat.
The win is the **bottleneck flip (GPU?, CPU?)**; the 2.9x is vs the untuned streaming path ? vs
prefetch-tuned cpu the realistic end-to-end is ~1.6?1.8x (#506 spike). All in the PR #507
comment.
47
48     ## Staged assets (durable)
49     - **Golden video working set**: all 14 originals in `gs://khelsutra-rally-
corpus/videos/golden/<name>.mp4` + the 4 baseline-source 1080p proxies in `videos/golden_1080p/`
(mahadevpura_2, Badminton_BXH_2, Boxhill_Doubles, GX020094). Frame-count-verified against each
golden trajectory. (GOLDEN_VIDEOS.md could record this GCS location in a follow-up ? kept out of
PR #507 to not muddy its scope.)
50     - Gate trajectories: `gs://khelsutra-rally-nightly/dev/506/results/golden/*_nvdec.csv`.
Re-score with repo-root `score_golden_gate.py` (untracked dev tool, like `score_local.py`;
labels from `?corpus/labels/`, baselines from `?corpus/trajectories/`).
51
52     ## Next steps
53     1. **Merge PR #507** ? real-L4-validated + golden-gate PASS; CI green locally.
54     2. **Flip the default**: owner Launch Report ? set `decode_backend=nvdec_resident` in
the cloud-serving preset (`backend/config/presets.py`). The gate that gated this is now green.
55
56     ## Ramp-up kit
57     - PR #507 (full change description + verification record) and issue #506 (all comments ?
measured evidence).
58     - New tests: `tests/test_wasb_nvdec_resident.py` (grid/cv2 parity incl. non-16:9, cache
keying, CLI guards, the skippable on-box parity gate).
59     - Reference code still in `scratch/gpu_resident/` (`exp_t2_full.py`, `score_local.py`,
`t2setup_env.sh`, `telemetry.sh`).
60
61     ## Key decisions (don't relitigate)
62     - GPU resize MUST be affine `grid_sample` matching `cv2.warpAffine` ? NOT
`F.interpolate` (F1 0.58?0.07; only shows on hard clips). Locked by unit tests now.
63     - `padding_mode="zeros"` (cv2's BORDER_CONSTANT(0)), not "border" ? matters only for
non-16:9 sources; identical to the validated run on 16:9.
64     - PyNvVideoCodec (MIT) for GPU decode; TorchCodec couldn't load on the DL image.
65     - Score with SERVED preset (cap6+R4); raw (0.02,2.0,2.0) is degenerate.
66     - Default-OFF; nvdec/cpu detection caches never cross-replay (manifest keyed by
decode_backend).
67     - WASB weights/NYCU-dataset license (WASB-C3) is a separate launch gate, not a perf
blocker.
68

```

5. user (2026-07-09T06:23:57.609Z)

<system-reminder>This memory is 2 days old. Memories are point-in-time observations, not live state ? claims about code behavior or file:line citations may be outdated. Verify against current code before asserting as fact.</system-reminder>

```

1     ---
2     name: efficient-timeouts-and-retries

```

```

3      description: How the owner wants Claude to handle waits, retries, timeouts, and Windows
gcloud SSH ? fast, visible, diagnostic
4      metadata:
5          type: feedback
6      ---
7
8      The owner strongly dislikes slow, silent retry loops. On 2026-07-06, diagnosing a GPU
VM, I burned ~3 minutes on a 12x15s SSH-readiness loop that suppressed stderr with 2>/dev/null,
so it "failed" invisibly while the owner had already SSH'd into the box manually. The owner
interrupted repeatedly ("be really smart at your job").
9
10     **Why:** wasteful polling plus hidden errors makes me look slow and not smart; the owner
watches actively and interrupts. One visible failure I actually diagnose beats many blind
retries.
11
12     **How to apply:**
13     - ONE attempt with visible stderr, read the real error, then fix ? do not blind-retry a
broken command.
14     - Short, reasonable timeouts (seconds, not minutes). Escalate only with a concrete
reason.
15     - Never hide stderr (2>/dev/null) on a command whose success is uncertain.
16     - Root cause: on Windows, `gcloud compute ssh` uses PuTTY plink.exe/pscp.exe, which
reject OpenSSH `-o` flags and mishandle some `~` remote paths. Use plain `gcloud compute ssh
<vm> --zone --project --command="..."` to push a file, base64-encode into --command (echo B64 |
base64 -d > file), not scp. To launch detached work without plink hanging ~60s, redirect all fds
and close stdin.
17     - For long remote work: run detached (nohup, log to file/GCS), poll a status file; avoid
many SSH round-trips.
18

```

6. assistant (2026-07-09T06:24:06.291Z)

The memory confirms the real code is at `C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-indexer`. Let me check its git state and pull per convention. Let me also start reading the target files in parallel.

7. user (2026-07-09T06:24:10.818Z)

```

=== branch ===
codex/issue-523-sd-reel-plan-clear
=== status ===
?? "SESSION_HANDOFF (1).md"
=== last commit ===
a460d74 chore(handoff): close issue #523 and move to next open bugs
=== fetch ===
=== behind/ahead vs origin/master ===
6      10

```

8. user (2026-07-09T06:24:18.429Z)

```

=== origin/master HEAD ===
00e8f92 feat(gates): local release perf gate + consolidate nightly into scripts/gates/ (#547)
2084d56 fix(storage): build GCSFuse fused path with OS-native separators (#548)
3249451 feat(#534): live detector progress heartbeat + L4 worker resource telemetry (#546)
=== does 00e8f92 match? ===
00e8f92

```

9. user (2026-07-09T06:24:18.455Z)

```

1      """"Worker dispatcher ? pull from storage ? run the pipeline ? push outputs back.
2
3      python -m backend.worker infer --video-uri <uri> --out-uri <prefix> \
4          [--segmenter wasb_hybrid] [--config config.json] [--set k=v ...] [--scratch DIR]
5      python -m backend.worker train ...          # Gen-0 train-from-bucket (follow-up)

```

```

6
7     The URI scheme selects the storage backend via ``backend.storage`` (a bare path /
``local://``
8     is an identity passthrough; ``s3://`` / ``gcs://`` is an object store), so existing
local runs
9     are byte-for-byte unchanged. Combine with ``RALLY_DETECTOR_IMPL=stub`` and
10    ``--set indexing.skip_ai_handoff=true`` for a **secret-free, GPU-free** run (no Gemini
key) ?
11    the path validated end-to-end on DigitalOcean (``docs/SETUP_NEW_MACHINE.md``).
12
13    This is a thin orchestrator: it composes the SAME building blocks the ``backend.main``
CLI uses
14    (validator registry ? segmenter ? report), never a forked pipeline.
15    """
16
17    from __future__ import annotations
18
19    import argparse
20    import csv
21    import json
22    import logging
23    import os
24    import re
25    import sys
26    import tempfile
27    from typing import Any, List
28
29    from backend import storage
30    from backend.compute import ComputeJobSpec, get_compute_backend
31    from backend.config import (
32        StartupCheckError,
33        enforce_startup_checks,
34        load_config,
35        probe_cuda_available,
36    )
37    from backend.infrastructure.database import Database
38    from backend.pipeline.segmenters.base import segmenter_registry
39    from backend.pipeline.validators.base import registry as validator_registry
40    from backend.reporting import emit_indexer_report
41    from backend.results import Failure
42    from backend.utils.ffmpeg import ffprobe_bin # P2a: single ffmpeg/ffprobe resolver
43    from backend.utils.hashing import compute_video_id
44
45    logger = logging.getLogger("backend.worker")
46
47
48    def _coerce(v: str) -> Any:
49        """Coerce a ``--set`` string value to bool/int/float where it round-trips, else
str."""
50        low = v.lower()
51        if low in ("true", "false"):
52            return low == "true"
53        for cast in (int, float):
54            try:
55                return cast(v)
56            except ValueError:
57                pass
58        return v
59
60
61    def _apply_overrides(config: Any, sets: List[str]) -> None:
62        """Apply ``a.b.c=value`` overrides onto the typed config (e.g.
indexing.skip_ai_handoff=true)."""
63        for kv in sets or []:
64            key, sep, val = kv.partition("=")

```

```

65         if not sep:
66             raise ValueError(f"--set expects k=v, got {kv!r}")
67         parts = key.split(".")
68         obj = config
69         for p in parts[:-1]:
70             obj = getattr(obj, p)
71         setattr(obj, parts[-1], _coerce(val))
72
73
74     def _write_intervals_csv(segments: List[dict], path: str) -> None:
75         """Decimal-second ``[start,end)`` rallies + shot count / ending reason ? the join-
friendly
76         artifact (same schema as the rally-annotator labels)."""
77         with open(path, "w", newline="") as f:
78             w = csv.writer(f)
79             w.writerow(["start", "end", "shot_count", "ending_reason"])
80             for s in segments:
81                 w.writerow(
82                     [
83                         f"{float(s.get('start_time', 0.0)):.3f}",
84                         f"{float(s.get('end_time', 0.0)):.3f}",
85                         s.get("shot_count", ""),
86                         s.get("ending_reason", s.get("shot_count_confidence", "")),
87                     ]
88                 )
89
90
91     def _write_report_json(
92         video_id: str, segments: List[dict], status: str, path: str, video_uri: str = ""
93     ) -> None:
94         with open(path, "w") as f:
95             json.dump(
96                 {
97                     "video_id": video_id,
98                     "status": status,
99                     "segment_count": len(segments),
100                     # The SOURCE the detection ran on (CP-rebuild P3, #485): a restart-wiped
101                     # control plane re-hydrating from outputs/<md5>/ restores the video row
102                     # with its truthful path ? without this, only the segments recover.
103                     "video_uri": video_uri,
104                     "segments": [
105                         {
106                             "start": float(s.get("start_time", 0.0)),
107                             "end": float(s.get("end_time", 0.0)),
108                         }
109                         for s in segments
110                     ],
111                 },
112                 f,
113                 indent=2,
114             )
115
116
117     def _serialize_and_push_outputs(
118         scratch: str,
119         out_uri: str,
120         video_id: str,
121         segments: List[dict],
122         status: str,
123         storage_cfg: dict,
124         video_uri: str = "",
125     ) -> List[str]:
126         """Write intervals.csv + report.json under ``<scratch>/outputs/<video_id>/`` and
push both to
127         ``<out_uri>/outputs/<video_id>/``; return the pushed URIs.

```



```

128
129     Shared by the SUCCESS path and the failure path (`_fail`) so a detector
timeout/abort leaves a
130     ``status="failed"`` report.json in the bucket ? the control plane's source of truth
131     (`orchestration.cached_process_result` / `probe_process_cache` treat any
non-``success``
132     report as "no usable result" ? retry, and the restart-recovery poll waits for
report.json to
133     EXIST) ? instead of a missing artifact or a false 0-segment ``"success"``.``"``
134     out_local = os.path.join(scratch, "outputs", video_id)
135     os.makedirs(out_local, exist_ok=True)
136     _write_intervals_csv(segments, os.path.join(out_local, "intervals.csv"))
137     _write_report_json(
138         video_id,
139         segments,
140         status,
141         os.path.join(out_local, "report.json"),
142         video_uri=video_uri,
143     )
144     out_prefix = out_uri.rstrip("/")
145     pushed: List[str] = []
146     for fn in sorted(os.listdir(out_local)):
147         uri = f"{out_prefix}/outputs/{video_id}/{fn}"
148         storage.put(os.path.join(out_local, fn), uri, config=storage_cfg)
149         pushed.append(uri)
150     return pushed
151
152
153 def _run_startup_checks(config: Any) -> bool:
154     """M5 per-profile startup checks for the worker (the compute process, so it probes
CUDA
155     best-effort). Logs warnings; returns False on a FATAL coherence error so the caller
aborts
156     cleanly (rc 2) rather than run an expensive job under an incoherent profile. A true
no-op
157     (returns True, ZERO work) when no deployment.profile is selected.
158
159     The CUDA probe is gated behind an active profile: `probe_cuda_available()` lazily
imports
160     torch, and pre-M5 the worker's infer/train path was torch-free (detection is
delegated to a
161     WSL/native subprocess). Probing only when a profile is set keeps the default-OFF
path a true
162     no-op of work, not just of output."""
163     profile = getattr(getattr(config, "deployment", None), "profile", "") or ""
164     cuda = (
165         probe_cuda_available() if profile else None
166     ) # torch-free on the default-OFF path
167     try:
168         enforce_startup_checks(config, cuda_available=cuda, logger=logger)
169         return True
170     except StartupCheckError:
171         return False # already logged at ERROR by enforce_startup_checks
172
173
174 def _dispatch_remote(compute: Any, args: argparse.Namespace) -> int:
175     """M6: hand ONE infer job to a remote ComputeBackend (Cloud Run) and return its rc
once the
176     outputs are durably in the bucket (the backend bucket-polls the done-marker). The
dispatched
177     container runs this SAME cmd_infer INLINE (compute_target forced local_gpu) ? never
re-dispatch.
178
179     A remote job that never writes a completion marker exits via one of two CLEAN codes
(distinct

```

```

180         from the local 2/3), never a raw traceback:
181         * rc 4 ? the bucket poll exhausted its budget (`PolicyTimeout`): the execution
may still be
182         running (the shim best-effort cancels it).
183         * rc 5 ? the execution has TERMINATED without a marker
(`RemoteExecutionFailed`): a worker
184         crash/OOM/pre-marker abort caught FAST by the Executions-API probe, so instead
of the
185         ~65-min opaque wait the operator gets an immediate, log-pointing failure."""
186     from backend.compute import RemoteExecutionFailed
187     from backend.infrastructure.execution_policy import PolicyTimeout
188
189     spec = ComputeJobSpec(
190         video_uri=args.video_uri,
191         out_uri=args.out_uri,
192         segmenter=args.segmenter,
193         config_overrides=tuple(args.set or ()),
194         skip_validation=bool(getattr(args, "skip_validation", False)),
195     )
196     try:
197         result = compute.run_infer(spec)
198     except RemoteExecutionFailed as e:
199         logger.error(
200             "[worker] remote execution terminated without a completion marker: %s", e
201         )
202         return 5
203     except PolicyTimeout as e:
204         logger.warning(
205             "[worker] remote compute did not complete (no marker within budget): %s", e
206         )
207         return 4
208     logger.info(
209         "[worker] remote compute done: video_id=%s rc=%d outputs=%s",
210         result.video_id,
211         result.returncode,
212         result.outputs_uri,
213     )
214     return result.returncode
215
216
217     def _write_done_marker(
218         done_uri: str,
219         video_id: str,
220         returncode: int,
221         segment_count: int,
222         status: str,
223         storage_cfg: dict,
224         scratch: str,
225     ) -> None:
226         """M6: write a tiny JSON completion marker to ``done_uri`` (via the storage layer)
for a remote
227         dispatcher's bucket-poll. Best-effort ? never raises into the worker's success
return."""
228         try:
229             marker = {
230                 "video_id": video_id,
231                 "returncode": returncode,
232                 "segment_count": segment_count,
233                 "status": status,
234             }
235             local = os.path.join(scratch, "_done.json")
236             with open(local, "w", encoding="utf-8") as f:
237                 json.dump(marker, f)
238             storage.put(local, done_uri, config=storage_cfg)
239             logger.info("[worker] wrote completion marker -> %s", done_uri)

```

```

240         except Exception as e: # never fail a good run because the marker push hiccuped
241             logger.warning("[worker] could not write done-marker %s: %s", done_uri, e)
242
243
244     def cmd_infer(args: argparse.Namespace) -> int:
245         config = load_config(args.config) if args.config else load_config()
246         _apply_overrides(config, args.set)
247         if not _run_startup_checks(config):
248             return 2
249
250         # M6: a cloud compute_target DISPATCHES to a Cloud Run Job instead of running in-
process.
251         # DEFAULT-OFF: get_compute_backend returns None for ""/local_gpu/cpu_mock ? fall
through to the
252         # unchanged inline path below. The dispatched container runs THIS cmd_infer inline
(the
253         # dispatcher forces compute_target=local_gpu), so it never recursively re-
dispatches.
254         compute = get_compute_backend(config)
255         if compute is not None:
256             return _dispatch_remote(compute, args)
257
258         storage_cfg = config.storage.model_dump()
259
260         scratch = args.scratch or tempfile.mkdtemp(prefix="rally-worker-")
261         os.makedirs(scratch, exist_ok=True)
262         config.output.output_dir = scratch
263
264         # 1. PULL ? local:// / bare path = identity passthrough; s3:// / gcs:// = download
to scratch.
265         local_video = storage.fetch(
266             args.video_uri, os.path.join(scratch, "in.mp4"), config=storage_cfg
267         )
268         print(f"[worker] pulled {args.video_uri} -> {local_video}")
269
270         # 2. IDENTITY ? whole-file MD5 (the labels?video join key; pulled bytes must be
byte-identical).
271         video_id = compute_video_id(local_video)
272
273         # 3. VALIDATE (same registry as the main CLI) ? UNLESS the dispatcher already did.
274         # On the cloud-serving path the CONTROL-PLANE is the authoritative validation gate:
it runs the
275         # SAME validators with ITS resolved config and only dispatches on PASS. Re-
validating here is
276         # redundant, decodes the whole video a SECOND time, and is outright WRONG when the
worker's config
277         # resolves different thresholds than the control-plane's ? observed 2026-07-07: a
5312x2988 clip
278         # (sharpness 11.9) PASSED the control-plane at validation.min_blur_threshold=8.0,
then the worker
279         # REJECTED it at its own default 20.0, so a SUCCESSFUL GPU dispatch failed after the
fact.
280         # --skip-validation (set by the dispatcher in orchestration._maybe_dispatch_remote)
skips it; the
281         # direct ``worker infer`` CLI has no control-plane, so it still validates by
default.
282         if getattr(args, "skip_validation", False):
283             print(
284                 "[worker] validation SKIPPED (--skip-validation): the dispatcher is the
authoritative "
285                 "gate and already validated this input."
286             )
287             val_results, val_context = [], {}
288         else:
289             val_results, val_context = validator_registry.run_all_with_context(

```

```

290         local_video, config
291     )
292     failures = [r for r in val_results if not r.passed]
293     db = Database(os.path.join(scratch, config.output.db_name))
294     db.add_video(
295         video_id,
296         local_video,
297         "failed" if failures else "processed",
298         [r.model_dump() for r in val_results],
299     )
300
301     def _fail(rc: int) -> int:
302         # A worker-side FAILURE (validation, unknown segmenter, or a detector
303         # segmenter surfaced as a Failure) must NOT masquerade as a 0-rally "success".
304         # a status="failed" report.json/intervals.csv (0 segments) so the control
305         # truth (submit-cache probe, restart-recovery) and the alpha-user result surface
306         # explicit FAILURE ? the wf294 silent-failure fix. Best-effort: a push hiccup
307         # mask the original failure. THEN the M6 done-marker (real rc + status) so a
308         # dispatcher's bucket-poll terminates FAST instead of hanging until its budget
309         # expires.
310         try:
311             _serialize_and_push_outputs(
312                 scratch,
313                 args.out_uri,
314                 video_id,
315                 [],
316                 "failed",
317                 storage_cfg,
318                 str(getattr(args, "video_uri", "") or ""),
319             )
320         except Exception as e: # never mask the real failure on an artifact-push hiccup
321             logger.warning("[worker] could not push failure artifacts: %s", e)
322             # DEFAULT-OFF: no --done-uri ? no marker written (the unchanged local run).
323             if getattr(args, "done_uri", None):
324                 _write_done_marker(
325                     args.done_uri, video_id, rc, 0, "failed", storage_cfg, scratch
326                 )
327             return rc
328
329     segments: List[dict] = []
330     segmenter_actual = None
331     segmentation_ran = False
332     status = "failed"
333     try:
334         if failures:
335             print(
336                 "[worker] validation FAILED: "
337                 + "; ".join(f"{r.validator_name}: {r.message}" for r in failures)
338             )
339             return _fail(2)
340         seg_name = args.segmenter or config.indexing.default_segmenter
341         segmenter_cls = segmenter_registry.get(seg_name)
342         if not segmenter_cls:
343             print(
344                 f"[worker] unknown segmenter: {seg_name!r} "
345                 f"(have {list(segmenter_registry.list_available())})"
346             )
347             return _fail(2)
348         segmenter_actual = seg_name

```

```

348         result = segmenter_cls(db, config).process_video(
349             video_id, local_video, max_frames=args.max_frames
350         )
351         segmentation_ran = True
352         if isinstance(result, Failure):
353             print(f"[worker] segmenter FAILED: {result.message}")
354             return _fail(3)
355         segments = result.value if hasattr(result, "value") else result
356         status = "success"
357         # Loud, never-silent diagnostic: a 0-segment success is almost always a
misconfig, not an
358         # empty match ? the cold-start failure mode (the substrate detector yielded no
usable
359         # trajectories). Explain it so a run is analysable from the log alone (re-run -v
for more).
360         if not segments:
361             detector_impl = (
362                 os.environ.get("RALLY_DETECTOR_IMPL")
363                 or getattr(config.indexing, "detector_impl", None)
364                 or "auto"
365             )
366             logger.warning(
367                 "infer produced 0 segments for video_id=%s (segmenter=%s,
detector_impl=%s). "
368                 "The detector substrate yielded no usable trajectories/windows ? check
the detector "
369                 "logs above; common causes: a missing ${AI_PROVIDER}_API_KEY with "
370                 "skip_ai_handoff=False (the segmenter bails BEFORE detection ? set the
key or "
371                 "indexing.skip_ai_handoff=true; cloud-serving now fails this at
startup), native "
372                 "runner UNAVAILABLE (weights/repo/WASB_PYTHON env), the WASB env
produced no "
373                 "detections, or the input resolution differs from the tuned proxy
resolution. "
374                 "Re-run with -v for the native command + frame counts.",
375                 video_id,
376                 segmenter_actual,
377                 detector_impl,
378             )
379         finally:
380             # Best-effort per-video observability report (never raises) ? parity with the
main CLI.
381             emit_indexer_report(
382                 db=db,
383                 video_id=video_id,
384                 video_path=local_video,
385                 config=config,
386                 val_results=val_results,
387                 val_context=val_context,
388                 segmenter_requested=args.segmenter,
389                 segmenter_actual=segmenter_actual,
390                 was_reingest=False,
391                 segmentation_ran=segmentation_ran,
392             )
393
394         # 5. SERIALIZE + 6. PUSH (outputs/<video_id>/?). idempotent on video_id.
395         pushed = _serialize_and_push_outputs(
396             scratch,
397             args.out_uri,
398             video_id,
399             segments,
400             status,
401             storage_cfg,
402             str(getattr(args, "video_uri", "") or ""),

```

```

403     )
404
405     print(
406         f"[worker] infer done: {len(segments)} segment(s); pushed {len(pushed)}
artifact(s):"
407     )
408     for u in pushed:
409         print("    ", u)
410
411     # M6: a remote dispatcher passes --done-uri; write the SUCCESS completion marker so
its bucket-poll
412     # terminates (carries video_id ? the dispatcher's output path + rc=0). Worker
FAILURES (rc 2/3)
413     # write their own marker via _fail() above, so the dispatcher fails FAST either way
(no 30-min hang
414     # on a bad video). DEFAULT-OFF: no --done-uri ? no marker written (the unchanged
local run).
415     if getattr(args, "done_uri", None):
416         _write_done_marker(
417             args.done_uri, video_id, 0, len(segments), status, storage_cfg, scratch
418         )
419     return 0
420
421
422     def _read_compile_spec(spec_uri: str, storage_cfg: dict, scratch: str) -> dict:
423         """Fetch + parse the control-plane-staged compile-spec (#521): the resolved
``time_pairs``, the
424         resolved ``stitching`` config, the source ``video_uri``, ``video_id``,
``output_name`` + ``locale``.
425         A gs://|s3:// URI is downloaded to scratch first (cloud backends require a real
dst)."""
426         local = storage.fetch(
427             spec_uri, os.path.join(scratch, "compile_spec.json"), config=storage_cfg
428         )
429         with open(local, encoding="utf-8") as f:
430             data = json.load(f)
431         if not isinstance(data, dict):
432             raise ValueError(f"compile spec at {spec_uri} is not a JSON object")
433         return data
434
435
436     def _write_compile_marker(
437         done_uri: str,
438         result: dict,
439         video_id: str,
440         returncode: int,
441         storage_cfg: dict,
442         scratch: str,
443     ) -> None:
444         """#521: write the compile done-marker (rc + status + the reel's
download_url/reel_uri) so the
445         dispatcher's bucket-poll terminates and can return the SPA result. Best-effort ?
never raises."""
446         try:
447             marker = {
448                 "video_id": video_id,
449                 "returncode": returncode,
450                 "status": result.get("status", "failed"),
451                 "download_url": result.get("download_url", ""),
452                 "reel_uri": result.get("reel_uri", ""),
453                 "filename": result.get("filename", ""),
454                 "rendition": result.get("rendition", ""),
455                 "message": result.get("message", ""),
456             }
457             local = os.path.join(scratch, "_compile_done.json")

```

```

458         with open(local, "w", encoding="utf-8") as f:
459             json.dump(marker, f)
460             storage.put(local, done_uri, config=storage_cfg)
461             logger.info("[worker] wrote compile marker -> %s", done_uri)
462     except Exception as e: # never fail a good stitch because the marker push hiccuped
463         logger.warning("[worker] could not write compile done-marker %s: %s", done_uri,
464             e)
465
466     def cmd_compile(args: argparse.Namespace) -> int:
467         """#521: run the reel COMPILE/stitch on THIS worker (the L4 has NVENC + 8 vCPU). The
468         stitch is the
469         2026-07-08 CUJ's dominant term (~13 min on the 2-vCPU control-plane for an 11-clip
470         4K reel). Read
471         the control-plane-staged compile-spec (resolved time-pairs + stitching config +
472         source URI), stitch
473         the reel via the SAME ``orchestration._stitch_reel`` core the control-plane uses (so
474         there is no
475         forked stitch pipeline), upload the reel to the bucket, and write a done-marker for
476         the dispatcher's
477         bucket-poll. No control-plane DB is needed ? the spec carries everything the stitch
478         requires."""
479         from backend import orchestration
480         from backend.config import StitchingConfig
481
482         config = load_config(args.config) if args.config else load_config()
483         _apply_overrides(config, args.set)
484         storage_cfg = config.storage.model_dump()
485
486         scratch = args.scratch or tempfile.mkdtemp(prefix="rally-compile-")
487         os.makedirs(scratch, exist_ok=True)
488         # Clips + the assembled reel land in scratch; the reel is then mirror-uploaded to
489         the bucket.
490         config.output.output_dir = scratch
491         # The dispatcher forwards storage.output_root via --set; --out-uri carries the same
492         value as a
493         # fallback so reel_object_uri (which reads storage.output_root) always resolves the
494         reel's bucket
495         # destination even if the override were dropped.
496         if not (config.storage.output_root or "").strip():
497             config.storage.output_root = args.out_uri
498
499         spec = _read_compile_spec(args.spec_uri, storage_cfg, scratch)
500         video_id = str(spec.get("video_id") or "")
501         out_name = str(spec.get("output_name") or "")
502         rendition = str(spec.get("rendition") or "original")
503         source_ref = str(spec.get("video_uri") or "")
504         locale = spec.get("locale")
505         time_pairs = [
506             (float(pair[0]), float(pair[1])) for pair in (spec.get("time_pairs") or [])
507         ]
508         stitching = StitchingConfig(**(spec.get("stitching") or {}))
509
510         print(
511             f"[worker] compile: video_id={video_id} clips={len(time_pairs)} "
512             f"encoder={stitching.encoder} rendition={rendition} out={out_name}"
513         )
514
515         result = orchestration._stitch_reel(
516             config,
517             stitching,
518             video_id,
519             source_ref,
520             time_pairs,
521             out_name,

```

```

513         locale,
514         lambda stage: print(f"[worker] compile: {stage}"),
515         rendition,
516         # On the L4 worker the bucket mirror is the ONLY reel delivery ? the worker's
local file is
517         # discarded when the Job exits. A lost mirror MUST fail the stitch (not a
404-on-"success").
518         require_mirror=True,
519     )
520     status = str(result.get("status") or "failed")
521     rc = 0 if status == "success" else 3
522     if status != "success":
523         print(f"[worker] compile FAILED: {result.get('message')}")
524
525     # DEFAULT-OFF: no --done-uri ? no marker (a direct-CLI stitch). A remote dispatcher
passes it.
526     if getattr(args, "done_uri", None):
527         _write_compile_marker(args.done_uri, result, video_id, rc, storage_cfg, scratch)
528     return rc
529
530
531     #: Video extensions accepted under <corpus>/videos/ (mirrors the labels/ '.csv' filter
so a
532     #: per-video sidecar ? .srt/.json/thumbnail ? can't shadow the real video on a stem
collision).
533     _VIDEO_EXTS = (".mp4", ".mov", ".mkv", ".avi", ".webm", ".m4v")
534
535     #: Leading ISO-date prefix on a canonical label name (`labels/<YYYY-MM-
DD>_<name>.rallies.csv`,
536     #: docs/DATA_IN_GCS.md §7.1). It disambiguates recycled GoPro names by authoring date
but is NOT
537     #: part of the clip's join key, so it's stripped when deriving the videos/ stem.
538     _LABEL_DATE_PREFIX = re.compile(r"^d{4}-d{2}-d{2}_")
539
540
541     def _label_clip_name(fname: str) -> str:
542         """Derive the clip name (the ``videos/<name>.<ext>`` join key) from a label
filename.
543
544         Strips both the ``.rallies.csv`` / ``.csv`` suffix and an optional leading ISO-date
prefix, so a
545         canonical dated label and a legacy bare label resolve to the SAME bare clip stem:
546
547         * ``2026-06-22_GX020094.rallies.csv`` -> ``GX020094`` (canonical, date-
stamped)
548         * ``GX020094.rallies.csv`` -> ``GX020094`` (legacy bare ? same
stem)
549         * ``GX020094_proxy.rallies.csv`` -> ``GX020094_proxy`` (``_proxy`` is part of
the name, kept)
550         * ``2026-06-21_mahadevpura_1.rallies.csv`` -> ``mahadevpura_1`` (underscores in
<name> survive)
551
552         Without the date strip, ``--trajectory-source detector`` can't match a dated label
to its bare
553         video stem, so every clip is skipped and train aborts with ``<2 videos``
(DATA_IN_GCS §7b #1).
554         """
555         name = fname.replace(".rallies.csv", "").replace(".csv", "")
556         return _LABEL_DATE_PREFIX.sub("", name)
557
558
559     def _probe_video_fps_width(path: str):
560         """Robust ffprobe (fps, frame_width) for a real video. Returns None on ANY failure.
561
562         The detector trajectory is frame-indexed; the harness maps frame->time via fps, so a

```



```

WRONG
563         fps silently mis-aligns every rally label (e.g. a 60fps clip read as 30 doubles
every time).
564         cv2's ``_probe_fps_width`` silently defaults to (30, 1280) on a read miss ? fine for
the
565         inference report (it flags the fallback) but corrupting for REAL training data. So
here we
566         probe with ffmpeg (same tool as ``get_video_duration``) and return None so the
caller SKIPS
567         rather than guesses.
568         """
569         import json
570         import subprocess
571
572         try:
573             out = subprocess.check_output(
574                 [
575                     ffmpeg_bin(),
576                     "-v",
577                     "error",
578                     "-select_streams",
579                     "v:0",
580                     "-show_entries",
581                     "stream=r_frame_rate,width",
582                     "-of",
583                     "json",
584                     path,
585                 ],
586                 stderr=subprocess.DEVNULL,
587             ).decode()
588             st = (json.loads(out).get("streams") or [{}])[0]
589             num, _, den = str(st.get("r_frame_rate", "0/1")).partition("/")
590             fps = float(num) / float(den or "1")
591             width = float(st.get("width") or 0)
592             if fps > 0 and width > 0:
593                 return fps, width
594         except (
595             subprocess.SubprocessError,
596             ValueError,
597             KeyError,
598             OSError,
599             json.JSONDecodeError,
600         ):
601             pass
602         return None
603
604
605     def _resolve_train_detector(args: argparse.Namespace, config: Any):
606         """Build the trajectory DetectorRunner for ``--trajectory-source detector``, honouring
607         detector_impl (RALLY_DETECTOR_IMPL / indexing.detector_impl: auto=WSL, native=GPU
box,
608         stub=CPU mock). Returns the runner, or prints an error + returns None.
609
610         Reuses the segmenter's ``_resolve_runner`` seam so the worker and the live CLI build
the
611         detector identically (one source of truth). A non-trajectory segmenter (e.g. yolo)
has
612         no shuttle detector and is rejected.
613         """
614         from backend.pipeline.segmenters.trajectory_hybrid import TrajectoryHybridSegmenter
615
616         seg_cls = segmenter_registry.get(args.segmenter)
617         if not seg_cls:
618             print(
619                 f"[worker] unknown segmenter: {args.segmenter!r} "

```

```

620         f"(have {list(segmenter_registry.list_available())})"
621     )
622     return None
623     seg = seg_cls(None, config)
624     if not isinstance(seg, TrajectoryHybridSegmenter):
625         print(
626             f"[worker] --segmenter {args.segmenter!r} has no shuttle detector; "
627             f"use a trajectory hybrid (wasb_hybrid / tracknet_hybrid / fusion_hybrid)."
628         )
629         return None
630     try:
631         runner, _ = seg._resolve_runner(config.indexing)
632     except NotImplementedError as e:
633         # e.g. detector_impl='native' for a family without a native runner (tracknet, or
634         # fusion with a tracknet substrate). Fail cleanly (rc!=0), not with a raw
635         traceback.print_exc()
636         print(f"[worker] detector not available: {e}")
637         return None
638     ok, msg = runner.healthcheck()
639     if not ok:
640         print(f"[worker] detector environment not ready: {msg}")
641         return None
642     print(f"[worker] detector ready ({args.segmenter}): {msg}")
643     return runner
644
645     def cmd_train(args: argparse.Namespace) -> int:
646         """Gen-0 train-from-bucket: pull labels (+ videos) -> trajectories -> manifest ->
647         shell out
648         to training.gen0.harness (backend must NOT import training) -> push the artifact
649         bundle.
650         Two trajectory sources (the train pipeline + harness are identical either way):
651         * ``synth`` (default) ? the CPU "mock GPU": deterministic, label-consistent
652         synthetic
653         shuttle paths (no GPU/WSL), so the whole pipeline is validated on a CPU box /
654         CI.
655         * ``detector`` ? REAL shuttle trajectories: pull each video and run the resolved
656         DetectorRunner (native WASB on a GPU box, or stub for a GPU-free wiring test).
657         This
658         is the M3.5 "first real training" path; only the trajectory source flips.
659         """
660         import subprocess
661
662         config = load_config(args.config) if args.config else load_config()
663         _apply_overrides(config, args.set)
664         if not _run_startup_checks(config):
665             return 2
666         storage_cfg = config.storage.model_dump()
667
668         scratch = args.scratch or tempfile.mkdtemp(prefix="rally-train-")
669         labels_dir = os.path.join(scratch, "labels")
670         traj_dir = os.path.join(scratch, "trajectories")
671         videos_dir = os.path.join(scratch, "videos")
672         os.makedirs(labels_dir, exist_ok=True)
673         os.makedirs(traj_dir, exist_ok=True)
674
675         corpus = args.corpus_uri.rstrip("/")
676         label_uris = sorted(
677             u
678             for u in storage.list_uris(f"{corpus}/labels/", config=storage_cfg)
679             if u.lower().endswith(".csv")
680         )
681         if len(label_uris) < 2:
682             print(

```

```

679         f"[worker] need >=2 labeled videos for leave-one-video-out; found
{len(label_uris)} "
680         f"under {corpus}/labels/"
681     )
682     return 2
683
684     # Real-detector source: resolve the runner once + index the bucket's videos/ by
stem.
685     runner = None
686     video_by_stem: dict = {}
687     if args.trajectory_source == "detector":
688         os.makedirs(videos_dir, exist_ok=True)
689         runner = _resolve_train_detector(args, config)
690         if runner is None:
691             return 2
692         for vuri in storage.list_uris(f"{corpus}/videos/", config=storage_cfg):
693             vname = storage.parse_uri(vuri).name
694             if not vname.lower().endswith(_VIDEO_EXTS):
695                 continue # skip sidecars (.srt/.json/thumbnails) so they can't shadow
the video
696             video_by_stem[os.path.splitext(vname)[0]] = vuri
697
698     print(f"[worker] trajectory source: {args.trajectory_source}")
699     specs: List[dict] = []
700     for uri in label_uris:
701         fname = storage.parse_uri(uri).name # e.g. 2026-06-22_GX020094.rallies.csv
702         name = _label_clip_name(fname) # -> GX020094 (date prefix stripped)
703         local_label = storage.fetch(
704             uri, os.path.join(labels_dir, fname), config=storage_cfg
705         )
706         if args.trajectory_source == "detector":
707             vuri = video_by_stem.get(name)
708             if not vuri:
709                 print(
710                     f"[worker] no video for {name!r} under {corpus}/videos/ ? skipping."
711                 )
712                 continue
713             local_video = storage.fetch(
714                 vuri,
715                 os.path.join(videos_dir, storage.parse_uri(vuri).name),
716                 config=storage_cfg,
717             )
718             video_id = compute_video_id(local_video)
719             # Real fps/width drive the trajectory frame->time mapping; PROBE (don't
guess) ? a
720             # wrong fps silently mis-aligns every label, so skip the video if ffprobe
can't read it.
721             meta = _probe_video_fps_width(local_video)
722             if meta is None:
723                 print(
724                     f"[worker] could not probe fps/width for {name!r} (ffprobe) ?
skipping (won't guess)."
725                 )
726                 continue
727             fps, frame_width = meta
728             traj = runner.run_predict(local_video, traj_dir, video_id=video_id)
729             if not traj:
730                 print(
731                     f"[worker] detector produced no trajectory for {name!r} ? skipping."
732                 )
733                 continue
734             specs.append(
735                 {
736                     "name": name,
737                     "trajectory": traj,

```

```

738         "labels": local_label,
739         "fps": fps,
740         "frame_width": frame_width,
741     }
742 )
743 else:
744     from backend.pipeline.detectors.stub_runner import (
745         synth_trajectory_from_labels,
746     )
747
748     traj = os.path.join(traj_dir, f"{name}_traj.csv")
749     synth_trajectory_from_labels(
750         local_label, traj, fps=args.fps, frame_width=args.frame_width
751     )
752     specs.append(
753         {
754             "name": name,
755             "trajectory": traj,
756             "labels": local_label,
757             "fps": args.fps,
758             "frame_width": args.frame_width,
759         }
760     )
761
762     if len(specs) < 2:
763         print(
764             f"[worker] need >=2 videos with trajectories for leave-one-video-out; got
{len(specs)}."
765         )
766         return 2
767     manifest_path = os.path.join(scratch, "manifest.json")
768     with open(manifest_path, "w", encoding="utf-8") as f:
769         json.dump(specs, f, indent=2)
770     src_label = (
771         "real detector" if args.trajectory_source == "detector" else "CPU-mock stub"
772     )
773     print(f"[worker] built manifest: {len(specs)} videos ({src_label} trajectories)")
774
775     out_dir = os.path.join(scratch, "artifacts")
776     cmd = [
777         sys.executable,
778         "-m",
779         "training.gen0.harness",
780         "--manifest",
781         manifest_path,
782         "--out",
783         out_dir,
784         "--version",
785         args.version,
786     ]
787     if args.floor:
788         floor_local = (
789             storage.fetch(
790                 args.floor, os.path.join(scratch, "floor.json"), config=storage_cfg
791             )
792             if "://" in args.floor
793             else args.floor
794         )
795         cmd += ["--floor", floor_local]
796     print("[worker] training:", " ".join(cmd))
797     rc = subprocess.run(cmd).returncode
798     if rc != 0:
799         print(f"[worker] gen0 harness failed (rc={rc})")
800         return rc
801

```

```

802     # PUSH the versioned artifact bundle (weights.json + manifest.json) to the bucket.
803     bundle = os.path.join(out_dir, args.version)
804     out_prefix = args.out_uri.rstrip("/")
805     pushed = []
806     for fn in sorted(os.listdir(bundle)):
807         uri = f"{out_prefix}/weights/{args.version}/{fn}"
808         storage.put(os.path.join(bundle, fn), uri, config=storage_cfg)
809         pushed.append(uri)
810     print(f"[worker] train done: pushed {len(pushed)} artifact(s):")
811     for u in pushed:
812         print("    ", u)
813     return 0
814
815
816 def build_parser() -> argparse.ArgumentParser:
817     p = argparse.ArgumentParser(
818         prog="python -m backend.worker",
819         description="Storage-aware worker: pull ? run pipeline ? push.",
820     )
821     p.add_argument(
822         "-v",
823         "--verbose",
824         action="store_true",
825         help="DEBUG logging (the native detector command, frame counts, per-stage
826 detail)",
827     )
828     sub = p.add_subparsers(dest="cmd", required=True)
829
830     pi = sub.add_parser("infer", help="run inference on one video URI and push outputs")
831     pi.add_argument(
832         "--video-uri", required=True, help="local path / local:// / s3:// / gcs://"
833     )
834     pi.add_argument(
835         "--out-uri",
836         required=True,
837         help="output prefix (outputs/<video_id>/? is appended)",
838     )
839     pi.add_argument(
840         "--segmenter", default=None, help="default: indexing.default_segmenter"
841     )
842     pi.add_argument(
843         "--config", default=None, help="config.json path (default: ./config.json)"
844     )
845     pi.add_argument(
846         "--scratch", default=None, help="local scratch dir (default: a temp dir)"
847     )
848     pi.add_argument("--max-frames", type=int, default=None)
849     pi.add_argument(
850         "--done-uri",
851         default=None,
852         help="M6: storage URI for a completion MARKER (video_id + rc) written on a "
853         "successful run, so a remote (Cloud Run) dispatcher's bucket-poll terminates. "
854         "Default-OFF: omit it for the normal local run.",
855     )
856     pi.add_argument(
857         "--set",
858         action="append",
859         default=[],
860         metavar="k=v",
861         help="config override, e.g. --set indexing.skip_ai_handoff=true",
862     )
863     pi.add_argument(
864         "--skip-validation",
865         action="store_true",
866         help="skip the worker's re-validation: the dispatcher (control-plane) already

```

```

ran the "
866         "validators with its resolved config and only dispatches on PASS, so re-
validating here is "
867         "redundant, decodes the video again, and can contradict the gate. Default-OFF:
the direct "
868         "`worker infer` CLI still validates.",
869     )
870     pi.set_defaults(func=cmd_infer)
871
872     pc = sub.add_parser(
873         "compile",
874         help="#521: stitch a highlight reel on this worker (L4 NVENC) from a staged
compile-spec",
875     )
876     pc.add_argument(
877         "--spec-uri",
878         required=True,
879         help="storage URI of the control-plane-staged compile-spec JSON (time-pairs +
stitching "
880         "config + source uri + video_id)",
881     )
882     pc.add_argument(
883         "--out-uri",
884         default="",
885         help="output ROOT (gs://? bucket); the reel is mirrored to
highlights/<video_id>/<name>. "
886         "Also forwarded via --set storage.output_root; either resolves the reel path.",
887     )
888     pc.add_argument(
889         "--done-uri",
890         default=None,
891         help="storage URI for the compile completion MARKER (rc + download_url), so a
remote "
892         "dispatcher's bucket-poll terminates. Default-OFF: omit for a direct-CLI
stitch.",
893     )
894     pc.add_argument(
895         "--config", default=None, help="config.json path (default: ./config.json)"
896     )
897     pc.add_argument(
898         "--scratch", default=None, help="local scratch dir (default: a temp dir)"
899     )
900     pc.add_argument(
901         "--set",
902         action="append",
903         default=[],
904         metavar="k=v",
905         help="config override, e.g. --set storage.output_root=gs://bucket",
906     )
907     pc.set_defaults(func=cmd_compile)
908
909     pt = sub.add_parser(
910         "train",
911         help="Gen-0 train-from-bucket (labels [+videos] -> trajectories -> harness)",
912     )
913     pt.add_argument(
914         "--corpus-uri",
915         required=True,
916         help="prefix containing labels/ (+ videos/ for "
917         "--trajectory-source detector), e.g. s3://bucket or gcs://bucket",
918     )
919     pt.add_argument(
920         "--out-uri",
921         required=True,
922         help="output prefix (weights/<version>/ is appended)",

```

```

923     )
924     pt.add_argument("--version", required=True, help="semver for the artifact bundle")
925     pt.add_argument(
926         "--trajectory-source",
927         choices=["synth", "detector"],
928         default="synth",
929         help="synth = CPU-mock label-consistent trajectories (default, no GPU); "
930         "detector = REAL shuttle trajectories via the resolved DetectorRunner "
931         "(needs videos/ in the corpus; honours detector_impl).",
932     )
933     pt.add_argument(
934         "--segmenter",
935         default="wasb_hybrid",
936         help="trajectory hybrid whose detector is used for --trajectory-source
detector",
937     )
938     pt.add_argument(
939         "--floor", default=None, help="floor baseline JSON (local path or storage URI)"
940     )
941     pt.add_argument("--config", default=None)
942     pt.add_argument("--scratch", default=None)
943     pt.add_argument(
944         "--fps",
945         type=float,
946         default=30.0,
947         help="synth trajectory fps (detector probes the video)",
948     )
949     pt.add_argument(
950         "--frame-width",
951         type=float,
952         default=1280.0,
953         help="synth trajectory frame width (detector probes the video)",
954     )
955     pt.add_argument("--set", action="append", default=[], metavar="k=v")
956     pt.set_defaults(func=cmd_train)
957     return p
958
959
960     def _setup_logging(verbose: bool) -> None:
961         """Configure logging for the worker CLI so the pipeline's existing INFO/ERROR
diagnostics are
962         VISIBLE ? the worker previously ran with no logging config, so every
segmenter/detector log
963         (detector choice, 'Running native WASB inference', availability failures) was
silently dropped.
964         That is exactly why a 0-rally cold-start gave nothing to analyse. Configured here
(the CLI
965         entry), not at import, so importing the module never reconfigures a caller's
logging."""
966         # force=True: basicConfig is a no-op if the root already has handlers (e.g. a
library
967         # configured logging at import), which would leave the level at WARNING and re-
suppress the
968         # diagnostics. Forcing reconfiguration means the worker reliably controls its own
verbosity.
969         logging.basicConfig(
970             level=logging.DEBUG if verbose else logging.INFO,
971             format="%(asctime)s [%(levelname)s] %(name)s: %(message)s",
972             datefmt="%H:%M:%S",
973             force=True,
974         )
975
976
977     def main(argv=None) -> int:
978         args = build_parser().parse_args(argv)

```

```

979         _setup_logging(getattr(args, "verbose", False))
980         return args.func(args)
981
982
983     if __name__ == "__main__":
984         sys.exit(main())
985

```

10. assistant (2026-07-09T06:24:49.365Z)

Origin/master is at `00e8f92` ? matches the task. Note the current branch `codex/issue-523-sd-reel-plan-clear` is ****6 behind**** origin/master, so it's missing PR #546 (the #534 heartbeat work) ? that's why the file I just read has no `\_progress` callback. The task's line numbers reference origin/master. Per branching-safety convention, let me cut my branch from `origin/master` explicitly and read the real target code there.

11. user (2026-07-09T06:24:52.289Z)

```

branch 'fix/worker-l4-infer-observability' set up to track 'origin/master'.
Switched to a new branch 'fix/worker-l4-infer-observability'
=== now on ===
fix/worker-l4-infer-observability
=== HEAD ===
00e8f92 feat(gates): local release perf gate + consolidate nightly into scripts/gates/ (#547)

```

12. assistant (2026-07-09T06:24:57.968Z)

Now on a clean branch off `origin/master`. Let me read the real target code (the origin/master version with the `\_progress` callback) and the two other files central to the fix.

13. user (2026-07-09T06:24:59.411Z)

```

1     """Worker dispatcher ? pull from storage ? run the pipeline ? push outputs back.
2
3     python -m backend.worker infer --video-uri <uri> --out-uri <prefix> \
4         [--segmenter wasb_hybrid] [--config config.json] [--set k=v ...] [--scratch DIR]
5     python -m backend.worker train ...          # Gen-0 train-from-bucket (follow-up)
6
7     The URI scheme selects the storage backend via ``backend.storage`` (a bare path /
8     ``local:`` / ``s3:`` / ``gcs:`` is an identity passthrough; ``s3:`` / ``gcs:`` is an object store), so existing
9     local runs
10    are byte-for-byte unchanged. Combine with ``RALLY_DETECTOR_IMPL=stub`` and
11    ``--set indexing.skip_ai_handoff=true`` for a **secret-free, GPU-free** run (no Gemini
12    key) ?
13    the path validated end-to-end on DigitalOcean (``docs/SETUP_NEW_MACHINE.md``).
14
15    This is a thin orchestrator: it composes the SAME building blocks the ``backend.main``
16    CLI uses
17    (validator registry ? segmenter ? report), never a forked pipeline.
18    """
19
20    from __future__ import annotations
21
22    import argparse
23    import csv
24    import json
25    import logging
26    import os
27    import re
28    import sys
29    import tempfile
30    import time
31    from typing import Any, List, Optional

```



```

29
30 from backend import storage
31 from backend.compute import ComputeJobSpec, get_compute_backend
32 from backend.config import (
33     StartupCheckError,
34     enforce_startup_checks,
35     load_config,
36     probe_cuda_available,
37 )
38 from backend.infrastructure.database import Database
39 from backend.pipeline.segmenters.base import segmenter_registry
40 from backend.pipeline.validators.base import registry as validator_registry
41 from backend.reporting import emit_indexer_report
42 from backend.results import Failure
43 from backend.utils.ffmpeg import ffprobe_bin # P2a: single ffmpeg/ffprobe resolver
44 from backend.utils.hashing import compute_video_id
45 from backend.utils.run_telemetry import RunTelemetry
46
47 logger = logging.getLogger("backend.worker")
48
49
50 def _coerce(v: str) -> Any:
51     """Coerce a ``--set`` string value to bool/int/float where it round-trips, else
52     str."""
53     low = v.lower()
54     if low in ("true", "false"):
55         return low == "true"
56     for cast in (int, float):
57         try:
58             return cast(v)
59         except ValueError:
60             pass
61     return v
62
63 def _apply_overrides(config: Any, sets: List[str]) -> None:
64     """Apply ``a.b.c=value`` overrides onto the typed config (e.g.
65     indexing.skip_ai_handoff=true).
66     Supports dotted paths into plain dict fields (e.g. storage.gcs.mount_root=...) for
67     #535."""
68     for kv in sets or []:
69         key, sep, val = kv.partition("=")
70         if not sep:
71             raise ValueError(f"--set expects k=v, got {kv!r}")
72         parts = key.split(".")
73         obj = config
74         for p in parts[:-1]:
75             obj = getattr(obj, p)
76         last = parts[-1]
77         coerced = _coerce(val)
78         if isinstance(obj, dict):
79             obj[last] = coerced
80         else:
81             setattr(obj, last, coerced)
82
83 def _write_intervals_csv(segments: List[dict], path: str) -> None:
84     """Decimal-second ``[start,end)`` rallies + shot count / ending reason ? the join-
85     friendly
86     artifact (same schema as the rally-annotator labels)."""
87     with open(path, "w", newline="") as f:
88         w = csv.writer(f)
89         w.writerow(["start", "end", "shot_count", "ending_reason"])
90         for s in segments:
91             w.writerow(

```

```

90         [
91             f"{float(s.get('start_time', 0.0)):.3f}",
92             f"{float(s.get('end_time', 0.0)):.3f}",
93             s.get("shot_count", ""),
94             s.get("ending_reason", s.get("shot_count_confidence", "")),
95         ]
96     )
97
98
99 def _write_report_json(
100     video_id: str, segments: List[dict], status: str, path: str, video_uri: str = ""
101 ) -> None:
102     with open(path, "w") as f:
103         json.dump(
104             {
105                 "video_id": video_id,
106                 "status": status,
107                 "segment_count": len(segments),
108                 # The SOURCE the detection ran on (CP-rebuild P3, #485): a restart-wiped
109                 # control plane re-hydrating from outputs/<md5>/ restores the video row
110                 # with its truthful path ? without this, only the segments recover.
111                 "video_uri": video_uri,
112                 "segments": [
113                     {
114                         "start": float(s.get("start_time", 0.0)),
115                         "end": float(s.get("end_time", 0.0)),
116                     }
117                     for s in segments
118                 ],
119             },
120             f,
121             indent=2,
122         )
123
124
125 def _serialize_and_push_outputs(
126     scratch: str,
127     out_uri: str,
128     video_id: str,
129     segments: List[dict],
130     status: str,
131     storage_cfg: dict,
132     video_uri: str = "",
133 ) -> List[str]:
134     """Write intervals.csv + report.json under ``<scratch>/outputs/<video_id>/`` and
135     push both to
136     ``<out_uri>/outputs/<video_id>/``; return the pushed URIs.
137     Shared by the SUCCESS path and the failure path (``_fail``) so a detector
138     timeout/abort leaves a
139     ``status="failed"`` report.json in the bucket ? the control plane's source of truth
140     (``orchestration.cached_process_result`` / ``probe_process_cache`` treat any
141     non-``success``
142     report as "no usable result" ? retry, and the restart-recovery poll waits for
143     report.json to
144     EXIST) ? instead of a missing artifact or a false 0-segment ``"success"``.
145     out_local = os.path.join(scratch, "outputs", video_id)
146     os.makedirs(out_local, exist_ok=True)
147     _write_intervals_csv(segments, os.path.join(out_local, "intervals.csv"))
148     _write_report_json(
149         video_id,
150         segments,
151         status,
152         os.path.join(out_local, "report.json"),
153         video_uri=video_uri,

```

```

151         )
152         out_prefix = out_uri.rstrip("/")
153         pushed: List[str] = []
154         for fn in sorted(os.listdir(out_local)):
155             uri = f"{out_prefix}/outputs/{video_id}/{fn}"
156             storage.put(os.path.join(out_local, fn), uri, config=storage_cfg)
157             pushed.append(uri)
158         return pushed
159
160
161     def _run_startup_checks(config: Any) -> bool:
162         """M5 per-profile startup checks for the worker (the compute process, so it probes
163         CUDA
164         best-effort). Logs warnings; returns False on a FATAL coherence error so the caller
165         aborts
166         cleanly (rc 2) rather than run an expensive job under an incoherent profile. A true
167         no-op
168         (returns True, ZERO work) when no deployment.profile is selected.
169
170         The CUDA probe is gated behind an active profile: ``probe_cuda_available()`` lazily
171         imports
172         torch, and pre-M5 the worker's infer/train path was torch-free (detection is
173         delegated to a
174         WSL/native subprocess). Probing only when a profile is set keeps the default-OFF
175         path a true
176         no-op of work, not just of output."""
177         profile = getattr(getattr(config, "deployment", None), "profile", "") or ""
178         cuda = (
179             probe_cuda_available() if profile else None
180         ) # torch-free on the default-OFF path
181         try:
182             enforce_startup_checks(config, cuda_available=cuda, logger=logger)
183             return True
184         except StartupCheckError:
185             return False # already logged at ERROR by enforce_startup_checks
186
187
188     def _dispatch_remote(compute: Any, args: argparse.Namespace) -> int:
189         """M6: hand ONE infer job to a remote ComputeBackend (Cloud Run) and return its rc
190         once the
191         outputs are durably in the bucket (the backend bucket-polls the done-marker). The
192         dispatched
193         container runs this SAME cmd_infer INLINE (compute_target forced local_gpu) ? never
194         re-dispatch.
195
196         A remote job that never writes a completion marker exits via one of two CLEAN codes
197         (distinct
198         from the local 2/3), never a raw traceback:
199         * rc 4 ? the bucket poll exhausted its budget (``PolicyTimeout``): the execution
200         may still be
201         running (the shim best-effort cancels it).
202         * rc 5 ? the execution has TERMINATED without a marker
203         (``RemoteExecutionFailed``): a worker
204         crash/OOM/pre-marker abort caught FAST by the Executions-API probe, so instead
205         of the
206         ~65-min opaque wait the operator gets an immediate, log-pointing failure."""
207         from backend.compute import RemoteExecutionFailed
208         from backend.infrastructure.execution_policy import PolicyTimeout
209
210         spec = ComputeJobSpec(
211             video_uri=args.video_uri,
212             out_uri=args.out_uri,
213             segmenter=args.segmenter,
214             config_overrides=tuple(args.set or ()),
215             skip_validation=bool(getattr(args, "skip_validation", False)),

```

```

203     )
204     try:
205         result = compute.run_infer(spec)
206     except RemoteExecutionFailed as e:
207         logger.error(
208             "[worker] remote execution terminated without a completion marker: %s", e
209         )
210         return 5
211     except PolicyTimeout as e:
212         logger.warning(
213             "[worker] remote compute did not complete (no marker within budget): %s", e
214         )
215         return 4
216     logger.info(
217         "[worker] remote compute done: video_id=%s rc=%d outputs=%s",
218         result.video_id,
219         result.returncode,
220         result.outputs_uri,
221     )
222     return result.returncode
223
224
225     def _write_done_marker(
226         done_uri: str,
227         video_id: str,
228         returncode: int,
229         segment_count: int,
230         status: str,
231         storage_cfg: dict,
232         scratch: str,
233     ) -> None:
234         """M6: write a tiny JSON completion marker to ``done_uri`` (via the storage layer)
235         for a remote
236         dispatcher's bucket-poll. Best-effort ? never raises into the worker's success
237         return."""
238         try:
239             marker = {
240                 "video_id": video_id,
241                 "returncode": returncode,
242                 "segment_count": segment_count,
243                 "status": status,
244             }
245             local = os.path.join(scratch, "_done.json")
246             with open(local, "w", encoding="utf-8") as f:
247                 json.dump(marker, f)
248             storage.put(local, done_uri, config=storage_cfg)
249             logger.info("[worker] wrote completion marker -> %s", done_uri)
250         except Exception as e: # never fail a good run because the marker push hiccuped
251             logger.warning("[worker] could not write done-marker %s: %s", done_uri, e)
252
253
254     def cmd_infer(args: argparse.Namespace) -> int:
255         config = load_config(args.config) if args.config else load_config()
256         _apply_overrides(config, args.set)
257         if not _run_startup_checks(config):
258             return 2
259
260         # M6: a cloud compute_target DISPATCHES to a Cloud Run Job instead of running in-
261         # process.
262         # DEFAULT-OFF: get_compute_backend returns None for ""/local_gpu/cpu_mock ? fall
263         # through to the
264         # unchanged inline path below. The dispatched container runs THIS cmd_infer inline
265         # (the
266         # dispatcher forces compute_target=local_gpu), so it never recursively re-
267         # dispatches.

```

```

262     compute = get_compute_backend(config)
263     if compute is not None:
264         return _dispatch_remote(compute, args)
265
266     storage_cfg = config.storage.model_dump()
267
268     scratch = args.scratch or tempfile.mkdtemp(prefix="rally-worker-")
269     os.makedirs(scratch, exist_ok=True)
270     config.output.output_dir = scratch
271
272     # 1. PULL ? local:// / bare path = identity passthrough; s3:// / gcs:// = download
to scratch.
273     local_video = storage.fetch(
274         args.video_uri, os.path.join(scratch, "in.mp4"), config=storage_cfg
275     )
276     print(f"[worker] pulled {args.video_uri} -> {local_video}")
277
278     # 2. IDENTITY ? whole-file MD5 (the labels?video join key; pulled bytes must be
byte-identical).
279     video_id = compute_video_id(local_video)
280
281     # 3. VALIDATE (same registry as the main CLI) ? UNLESS the dispatcher already did.
282     # On the cloud-serving path the CONTROL-PLANE is the authoritative validation gate:
it runs the
283     # SAME validators with ITS resolved config and only dispatches on PASS. Re-
validating here is
284     # redundant, decodes the whole video a SECOND time, and is outright WRONG when the
worker's config
285     # resolves different thresholds than the control-plane's ? observed 2026-07-07: a
5312x2988 clip
286     # (sharpness 11.9) PASSED the control-plane at validation.min_blur_threshold=8.0,
then the worker
287     # REJECTED it at its own default 20.0, so a SUCCESSFUL GPU dispatch failed after the
fact.
288     # --skip-validation (set by the dispatcher in orchestration._maybe_dispatch_remote)
skips it; the
289     # direct ``worker infer`` CLI has no control-plane, so it still validates by
default.
290     if getattr(args, "skip_validation", False):
291         print(
292             "[worker] validation SKIPPED (--skip-validation): the dispatcher is the
authoritative "
293             "gate and already validated this input."
294         )
295         val_results, val_context = [], {}
296     else:
297         val_results, val_context = validator_registry.run_all_with_context(
298             local_video, config
299         )
300     failures = [r for r in val_results if not r.passed]
301     db = Database(os.path.join(scratch, config.output.db_name))
302     db.add_video(
303         video_id,
304         local_video,
305         "failed" if failures else "processed",
306         [r.model_dump() for r in val_results],
307     )
308
309     def _fail(rc: int) -> int:
310         # A worker-side FAILURE (validation, unknown segmenter, or a detector
timeout/abort the
311         # segmenter surfaced as a Failure) must NOT masquerade as a 0-rally "success".
Write + push
312         # a status="failed" report.json/intervals.csv (0 segments) so the control
plane's bucket

```

```

313         # truth (submit-cache probe, restart-recovery) and the alpha-user result surface
see an
314         # explicit FAILURE ? the wf294 silent-failure fix. Best-effort: a push hiccup
must never
315         # mask the original failure. THEN the M6 done-marker (real rc + status) so a
remote
316         # dispatcher's bucket-poll terminates FAST instead of hanging until its budget
expires.
317         try:
318             _serialize_and_push_outputs(
319                 scratch,
320                 args.out_uri,
321                 video_id,
322                 [],
323                 "failed",
324                 storage_cfg,
325                 str(getattr(args, "video_uri", "") or ""),
326             )
327         except Exception as e: # never mask the real failure on an artifact-push hiccup
328             logger.warning("[worker] could not push failure artifacts: %s", e)
329         # DEFAULT-OFF: no --done-uri ? no marker written (the unchanged local run).
330         if getattr(args, "done_uri", None):
331             _write_done_marker(
332                 args.done_uri, video_id, rc, 0, "failed", storage_cfg, scratch
333             )
334         return rc
335
336     segments: List[dict] = []
337     segmenter_actual = None
338     segmentation_ran = False
339     status = "failed"
340     try:
341         if failures:
342             print(
343                 "[worker] validation FAILED: "
344                 + "; ".join(f"{r.validator_name}: {r.message}" for r in failures)
345             )
346             return _fail(2)
347         seg_name = args.segmenter or config.indexing.default_segmenter
348         segmenter_cls = segmenter_registry.get(seg_name)
349         if not segmenter_cls:
350             print(
351                 f"[worker] unknown segmenter: {seg_name!r} "
352                 f"(have {list(segmenter_registry.list_available())})"
353             )
354             return _fail(2)
355         segmenter_actual = seg_name
356         out_local = os.path.join(scratch, "outputs", video_id) # for sidecar +
telemetry
357         os.makedirs(out_local, exist_ok=True)
358
359         # Live telemetry black-box for L4 detector runs (#534): velocity + GPU/CPU/RAM
samples.
360         # Written under outputs/ so the final push will mirror it; snapshots driven from
detector %.
361         telem: Optional[RunTelemetry] = None
362         try:
363             telem = RunTelemetry(
364                 os.path.join(out_local, "run_telemetry.jsonl"),
365                 label=f"l4-infer-{video_id[:8]}",
366             )
367         except Exception: # noqa: BLE001 - telemetry must never break a production run
368             telem = None
369
370         out_prefix = getattr(args, "out_uri", None)

```

```

371         if out_prefix:
372             out_prefix = out_prefix.rstrip("/")
373
374         # Configurable sampling for sidecar upload + telemetry snapshots (#534).
375         # Local sidecar file is always kept up-to-date (cheap). Only the costly remote
376         # and nvidia-smi-using snapshot are rate-limited. 0/negative = every callback.
377         sample_sec = 15.0
378         try:
379             idx_cfg = getattr(config, "indexing", None)
380             if idx_cfg is not None:
381                 raw = getattr(idx_cfg, "detector_progress_sample_sec", 15.0)
382                 sample_sec = float(raw) if raw is not None else 15.0
383         except Exception: # noqa: BLE001
384             sample_sec = 15.0
385
386         last_sample_t = [0.0]
387
388         def _progress(s: str) -> None:
389             print(f"[worker] progress: {s}")
390             # Write sidecar for cloud/L4 poll to surface live detector progress in job
391             # CP will read outputs/<video_id>/_detector_progress.json during poll.
392             # The *local* file is always written so the latest state is available in
393             # scratch.
394             now = time.time()
395             do_sample = sample_sec <= 0 or (now - last_sample_t[0] >= sample_sec)
396
397             pfile = os.path.join(out_local, f"{video_id}_detector_progress.json")
398             try:
399                 with open(pfile, "w") as pf:
400                     json.dump({"msg": s, "t": now}, pf)
401             except Exception:
402                 pass # best effort
403
404             heavy_done = False
405             # LIVE push to storage (the part visible to CP poll during the run).
406             if out_prefix and do_sample:
407                 try:
408                     storage.put(
409                         pfile,
410                         f"{out_prefix}/outputs/{video_id}/{os.path.basename(pfile)}",
411                         config=storage_cfg,
412                     )
413                     heavy_done = True
414                 except Exception:
415                     pass # best-effort live heartbeat
416
417             # Drive RunTelemetry snapshots (includes nvidia-smi + file rewrite) only on
418             # sample.
419             if telem and s and "detector" in (s or "") and do_sample:
420                 try:
421                     m = re.search(r"(\d+)/(\d+)", s)
422                     done = int(m.group(1)) if m else None
423                     total = int(m.group(2)) if m else None
424                     telem.snapshot(done=done, total=total, phase="detector")
425                     heavy_done = True
426                 except Exception: # noqa: BLE001
427                     pass
428
429             if heavy_done:
430                 last_sample_t[0] = now
431
432         result = segmenter_cls(db, config).process_video(
433             video_id,

```

```

432         local_video,
433         max_frames=args.max_frames,
434         progress=_progress,
435     )
436     segmentation_ran = True
437     if isinstance(result, Failure):
438         print(f"[worker] segmenter FAILED: {result.message}")
439         return _fail(3)
440     segments = result.value if hasattr(result, "value") else result
441     status = "success"
442     # Loud, never-silent diagnostic: a 0-segment success is almost always a
misconfig, not an
443     # empty match ? the cold-start failure mode (the substrate detector yielded no
usable
444     # trajectories). Explain it so a run is analysable from the log alone (re-run -v
for more).
445     if not segments:
446         detector_impl = (
447             os.environ.get("RALLY_DETECTOR_IMPL")
448             or getattr(config.indexing, "detector_impl", None)
449             or "auto"
450         )
451         logger.warning(
452             "infer produced 0 segments for video_id=%s (segmenter=%s,
detector_impl=%s). "
453             "The detector substrate yielded no usable trajectories/windows ? check
the detector "
454             "logs above; common causes: a missing ${AI_PROVIDER}_API_KEY with "
455             "skip_ai_handoff=False (the segmenter bails BEFORE detection ? set the
key or "
456             "indexing.skip_ai_handoff=true; cloud-serving now fails this at
startup), native "
457             "runner UNAVAILABLE (weights/repo/WASB_PYTHON env), the WASB env
produced no "
458             "detections, or the input resolution differs from the tuned proxy
resolution. "
459             "Re-run with -v for the native command + frame counts.",
460             video_id,
461             segmenter_actual,
462             detector_impl,
463         )
464     finally:
465         # Best-effort per-video observability report (never raises) ? parity with the
main CLI.
466         emit_indexer_report(
467             db=db,
468             video_id=video_id,
469             video_path=local_video,
470             config=config,
471             val_results=val_results,
472             val_context=val_context,
473             segmenter_requested=args.segmenter,
474             segmenter_actual=segmenter_actual,
475             was_reingest=False,
476             segmentation_ran=segmentation_ran,
477         )
478
479     # 5. SERIALIZE + 6. PUSH (outputs/<video_id>/?). idempotent on video_id.
480     pushed = _serialize_and_push_outputs(
481         scratch,
482         args.out_uri,
483         video_id,
484         segments,
485         status,
486         storage_cfg,

```



```

487         str(getattr(args, "video_uri", "") or ""),
488     )
489
490     print(
491         f"[worker] infer done: {len(segments)} segment(s); pushed {len(pushed)}"
492     )
493     for u in pushed:
494         print("    ", u)
495
496     # M6: a remote dispatcher passes --done-uri; write the SUCCESS completion marker so
497     # its bucket-poll
498     # terminates (carries video_id ? the dispatcher's output path + rc=0). Worker
499     FAILURES (rc 2/3)
500     # write their own marker via _fail() above, so the dispatcher fails FAST either way
501     # (no 30-min hang
502     # on a bad video). DEFAULT-OFF: no --done-uri ? no marker written (the unchanged
503     local run).
504     if getattr(args, "done_uri", None):
505         _write_done_marker(
506             args.done_uri, video_id, 0, len(segments), status, storage_cfg, scratch
507         )
508     return 0
509
510     def _read_compile_spec(spec_uri: str, storage_cfg: dict, scratch: str) -> dict:
511         """Fetch + parse the control-plane-staged compile-spec (#521): the resolved
512         ``time_pairs``, the
513         resolved ``stitching`` config, the source ``video_uri``, ``video_id``,
514         ``output_name`` + ``locale``.
515         A gs://|s3:// URI is downloaded to scratch first (cloud backends require a real
516         dst)."""
517         local = storage.fetch(
518             spec_uri, os.path.join(scratch, "compile_spec.json"), config=storage_cfg
519         )
520         with open(local, encoding="utf-8") as f:
521             data = json.load(f)
522         if not isinstance(data, dict):
523             raise ValueError(f"compile spec at {spec_uri} is not a JSON object")
524         return data
525
526     def _write_compile_marker(
527         done_uri: str,
528         result: dict,
529         video_id: str,
530         returncode: int,
531         storage_cfg: dict,
532         scratch: str,
533     ) -> None:
534         """#521: write the compile done-marker (rc + status + the reel's
535         download_url/reel_uri) so the
536         dispatcher's bucket-poll terminates and can return the SPA result. Best-effort ?
537         never raises."""
538         try:
539             marker = {
540                 "video_id": video_id,
541                 "returncode": returncode,
542                 "status": result.get("status", "failed"),
543                 "download_url": result.get("download_url", ""),
544                 "reel_uri": result.get("reel_uri", ""),
545                 "filename": result.get("filename", ""),
546                 "rendition": result.get("rendition", ""),
547                 "message": result.get("message", ""),
548             }

```

```

542         local = os.path.join(scratch, "_compile_done.json")
543         with open(local, "w", encoding="utf-8") as f:
544             json.dump(marker, f)
545             storage.put(local, done_uri, config=storage_cfg)
546             logger.info("[worker] wrote compile marker -> %s", done_uri)
547     except Exception as e: # never fail a good stitch because the marker push hiccuped
548         logger.warning("[worker] could not write compile done-marker %s: %s", done_uri,
549 e)
550
551     def cmd_compile(args: argparse.Namespace) -> int:
552         """#521: run the reel COMPILER/stitch on THIS worker (the L4 has NVENC + 8 vCPU). The
553 stitch is the
554         2026-07-08 CUJ's dominant term (~13 min on the 2-vCPU control-plane for an 11-clip
555 4K reel). Read
556         the control-plane-staged compile-spec (resolved time-pairs + stitching config +
557 source URI), stitch
558         the reel via the SAME ``orchestration._stitch_reel`` core the control-plane uses (so
559 there is no
560         forked stitch pipeline), upload the reel to the bucket, and write a done-marker for
561 the dispatcher's
562         bucket-poll. No control-plane DB is needed ? the spec carries everything the stitch
563 requires."""
564         from backend import orchestration
565         from backend.config import StitchingConfig
566
567         config = load_config(args.config) if args.config else load_config()
568         _apply_overrides(config, args.set)
569         storage_cfg = config.storage.model_dump()
570
571         scratch = args.scratch or tempfile.mkdtemp(prefix="rally-compile-")
572         os.makedirs(scratch, exist_ok=True)
573         # Clips + the assembled reel land in scratch; the reel is then mirror-uploaded to
574 the bucket.
575         config.output.output_dir = scratch
576         # The dispatcher forwards storage.output_root via --set; --out-uri carries the same
577 value as a
578         # fallback so reel_object_uri (which reads storage.output_root) always resolves the
579 reel's bucket
580         # destination even if the override were dropped.
581         if not (config.storage.output_root or "").strip():
582             config.storage.output_root = args.out_uri
583
584         spec = _read_compile_spec(args.spec_uri, storage_cfg, scratch)
585         video_id = str(spec.get("video_id") or "")
586         out_name = str(spec.get("output_name") or "")
587         rendition = str(spec.get("rendition") or "original")
588         source_ref = str(spec.get("video_uri") or "")
589         locale = spec.get("locale")
590         time_pairs = [
591             (float(pair[0]), float(pair[1])) for pair in (spec.get("time_pairs") or [])
592         ]
593         stitching = StitchingConfig(**(spec.get("stitching") or {}))
594
595         print(
596             f"[worker] compile: video_id={video_id} clips={len(time_pairs)} "
597             f"encoder={stitching.encoder} rendition={rendition} out={out_name} "
598             f"ranges={time_pairs}"
599         )
600
601         result = orchestration._stitch_reel(
602             config,
603             stitching,
604             video_id,
605             source_ref,

```

```

597         time_pairs,
598         out_name,
599         locale,
600         lambda stage: print(f"[worker] compile: {stage}"),
601         rendition,
602         # On the L4 worker the bucket mirror is the ONLY reel delivery ? the worker's
local file is
603         # discarded when the Job exits. A lost mirror MUST fail the stitch (not a
404-on-"success").
604         require_mirror=True,
605     )
606     status = str(result.get("status") or "failed")
607     rc = 0 if status == "success" else 3
608     if status != "success":
609         print(f"[worker] compile FAILED: {result.get('message')}")
610
611     # DEFAULT-OFF: no --done-uri ? no marker (a direct-CLI stitch). A remote dispatcher
passes it.
612     if getattr(args, "done_uri", None):
613         _write_compile_marker(args.done_uri, result, video_id, rc, storage_cfg, scratch)
614     return rc
615
616
617     #: Video extensions accepted under <corpus>/videos/ (mirrors the labels/ '.csv' filter
so a
618     #: per-video sidecar ? .srt/.json/thumbnail ? can't shadow the real video on a stem
collision).
619     _VIDEO_EXTS = (".mp4", ".mov", ".mkv", ".avi", ".webm", ".m4v")
620
621     #: Leading ISO-date prefix on a canonical label name (`labels/<YYYY-MM-
DD>_<name>.rallies.csv`,
622     #: docs/DATA_IN_GCS.md §7.1). It disambiguates recycled GoPro names by authoring date
but is NOT
623     #: part of the clip's join key, so it's stripped when deriving the videos/ stem.
624     _LABEL_DATE_PREFIX = re.compile(r"^\d{4}-\d{2}-\d{2}_")
625
626
627     def _label_clip_name(fname: str) -> str:
628         """Derive the clip name (the `videos/<name>.<ext>` join key) from a label
filename.
629
630         Strips both the `.rallies.csv` / `.csv` suffix and an optional leading ISO-date
prefix, so a
631         canonical dated label and a legacy bare label resolve to the SAME bare clip stem:
632
633         * `2026-06-22_GX020094.rallies.csv` -> `GX020094` (canonical, date-
stamped)
634         * `GX020094.rallies.csv` -> `GX020094` (legacy bare ? same
stem)
635         * `GX020094_proxy.rallies.csv` -> `GX020094_proxy` (`_proxy` is part of
the name, kept)
636         * `2026-06-21_mahadevpura_1.rallies.csv` -> `mahadevpura_1` (underscores in
<name> survive)
637
638         Without the date strip, `--trajectory-source detector` can't match a dated label
to its bare
639         video stem, so every clip is skipped and train aborts with `<2 videos`
(DATA_IN_GCS §7b #1).
640         """
641         name = fname.replace(".rallies.csv", "").replace(".csv", "")
642         return _LABEL_DATE_PREFIX.sub("", name)
643
644
645     def _probe_video_fps_width(path: str):
646         """Robust ffprobe (fps, frame_width) for a real video. Returns None on ANY failure.

```

```

647
648     The detector trajectory is frame-indexed; the harness maps frame->time via fps, so a
WRONG
649     fps silently mis-aligns every rally label (e.g. a 60fps clip read as 30 doubles
every time).
650     cv2's ``_probe_fps_width`` silently defaults to (30, 1280) on a read miss ? fine for
the
651     inference report (it flags the fallback) but corrupting for REAL training data. So
here we
652     probe with ffprobe (same tool as ``get_video_duration``) and return None so the
caller SKIPS
653     rather than guesses.
654     """
655     import json
656     import subprocess
657
658     try:
659         out = subprocess.check_output(
660             [
661                 ffprobe_bin(),
662                 "-v",
663                 "error",
664                 "-select_streams",
665                 "v:0",
666                 "-show_entries",
667                 "stream=r_frame_rate,width",
668                 "-of",
669                 "json",
670                 path,
671             ],
672             stderr=subprocess.DEVNULL,
673             ).decode()
674         st = (json.loads(out).get("streams") or [{}])[0]
675         num, _, den = str(st.get("r_frame_rate", "0/1")).partition("/")
676         fps = float(num) / float(den or "1")
677         width = float(st.get("width") or 0)
678         if fps > 0 and width > 0:
679             return fps, width
680     except (
681         subprocess.SubprocessError,
682         ValueError,
683         KeyError,
684         OSError,
685         json.JSONDecodeError,
686     ):
687         pass
688     return None
689
690
691 def _resolve_train_detector(args: argparse.Namespace, config: Any):
692     """Build the trajectory DetectorRunner for ``--trajectory-source detector``, honouring
693     detector_impl (RALLY_DETECTOR_IMPL / indexing.detector_impl: auto=WSL, native=GPU
box,
694     stub=CPU mock). Returns the runner, or prints an error + returns None.
695
696     Reuses the segmenter's ``_resolve_runner`` seam so the worker and the live CLI build
the
697     detector identically (one source of truth). A non-trajectory segmenter (e.g. yolo)
has
698     no shuttle detector and is rejected.
699     """
700     from backend.pipeline.segmenters.trajectory_hybrid import TrajectoryHybridSegmenter
701
702     seg_cls = segmenter_registry.get(args.segmenter)
703     if not seg_cls:

```

```

704         print(
705             f"[worker] unknown segmenter: {args.segmenter!r} "
706             f"(have {list(segmenter_registry.list_available())})"
707         )
708         return None
709     seg = seg_cls(None, config)
710     if not isinstance(seg, TrajectoryHybridSegmenter):
711         print(
712             f"[worker] --segmenter {args.segmenter!r} has no shuttle detector; "
713             f"use a trajectory hybrid (wasb_hybrid / tracknet_hybrid / fusion_hybrid)."
714         )
715         return None
716     try:
717         runner, _ = seg._resolve_runner(config.indexing)
718     except NotImplementedError as e:
719         # e.g. detector_impl='native' for a family without a native runner (tracknet, or
720         # fusion with a tracknet substrate). Fail cleanly (rc!=0), not with a raw
721         traceback.print_exc()
722         print(f"[worker] detector not available: {e}")
723         return None
724     ok, msg = runner.healthcheck()
725     if not ok:
726         print(f"[worker] detector environment not ready: {msg}")
727         return None
728     print(f"[worker] detector ready ({args.segmenter}): {msg}")
729     return runner
730
731 def cmd_train(args: argparse.Namespace) -> int:
732     """Gen-0 train-from-bucket: pull labels (+ videos) -> trajectories -> manifest ->
733     shell out
734     to training.gen0.harness (backend must NOT import training) -> push the artifact
735     bundle.
736     Two trajectory sources (the train pipeline + harness are identical either way):
737     * ``synth`` (default) ? the CPU "mock GPU": deterministic, label-consistent
738     synthetic
739     shuttle paths (no GPU/WSL), so the whole pipeline is validated on a CPU box /
740     CI.
741     * ``detector`` ? REAL shuttle trajectories: pull each video and run the resolved
742     DetectorRunner (native WASB on a GPU box, or stub for a GPU-free wiring test).
743     This
744     is the M3.5 "first real training" path; only the trajectory source flips.
745     """
746     import subprocess
747
748     config = load_config(args.config) if args.config else load_config()
749     _apply_overrides(config, args.set)
750     if not _run_startup_checks(config):
751         return 2
752     storage_cfg = config.storage.model_dump()
753
754     scratch = args.scratch or tempfile.mkdtemp(prefix="rally-train-")
755     labels_dir = os.path.join(scratch, "labels")
756     traj_dir = os.path.join(scratch, "trajectories")
757     videos_dir = os.path.join(scratch, "videos")
758     os.makedirs(labels_dir, exist_ok=True)
759     os.makedirs(traj_dir, exist_ok=True)
760
761     corpus = args.corpus_uri.rstrip("/")
762     label_uris = sorted(
763         u
764         for u in storage.list_uris(f"{corpus}/labels/", config=storage_cfg)
765         if u.lower().endswith(".csv")
766     )

```

```

763         if len(label_uris) < 2:
764             print(
765                 f"[worker] need >=2 labeled videos for leave-one-video-out; found
{len(label_uris)} "
766                 f"under {corpus}/labels/"
767             )
768             return 2
769
770     # Real-detector source: resolve the runner once + index the bucket's videos/ by
stem.
771     runner = None
772     video_by_stem: dict = {}
773     if args.trajectory_source == "detector":
774         os.makedirs(videos_dir, exist_ok=True)
775         runner = _resolve_train_detector(args, config)
776         if runner is None:
777             return 2
778         for vuri in storage.list_uris(f"{corpus}/videos/", config=storage_cfg):
779             vname = storage.parse_uri(vuri).name
780             if not vname.lower().endswith(_VIDEO_EXTS):
781                 continue # skip sidecars (.srt/.json/thumbnails) so they can't shadow
the video
782                 video_by_stem[os.path.splitext(vname)[0]] = vuri
783
784     print(f"[worker] trajectory source: {args.trajectory_source}")
785     specs: List[dict] = []
786     for uri in label_uris:
787         fname = storage.parse_uri(uri).name # e.g. 2026-06-22_GX020094.rallies.csv
788         name = _label_clip_name(fname) # -> GX020094 (date prefix stripped)
789         local_label = storage.fetch(
790             uri, os.path.join(labels_dir, fname), config=storage_cfg
791         )
792         if args.trajectory_source == "detector":
793             vuri = video_by_stem.get(name)
794             if not vuri:
795                 print(
796                     f"[worker] no video for {name!r} under {corpus}/videos/ ? skipping."
797                 )
798                 continue
799             local_video = storage.fetch(
800                 vuri,
801                 os.path.join(videos_dir, storage.parse_uri(vuri).name),
802                 config=storage_cfg,
803             )
804             video_id = compute_video_id(local_video)
805             # Real fps/width drive the trajectory frame->time mapping; PROBE (don't
guess) ? a
806             # wrong fps silently mis-aligns every label, so skip the video if ffprobe
can't read it.
807             meta = _probe_video_fps_width(local_video)
808             if meta is None:
809                 print(
810                     f"[worker] could not probe fps/width for {name!r} (ffprobe) ?
skipping (won't guess)."
811                 )
812                 continue
813             fps, frame_width = meta
814             traj = runner.run_predict(local_video, traj_dir, video_id=video_id)
815             if not traj:
816                 print(
817                     f"[worker] detector produced no trajectory for {name!r} ? skipping."
818                 )
819                 continue
820             specs.append(
821                 {

```

```

822         "name": name,
823         "trajectory": traj,
824         "labels": local_label,
825         "fps": fps,
826         "frame_width": frame_width,
827     }
828 )
829 else:
830     from backend.pipeline.detectors.stub_runner import (
831         synth_trajectory_from_labels,
832     )
833
834     traj = os.path.join(traj_dir, f"{name}_traj.csv")
835     synth_trajectory_from_labels(
836         local_label, traj, fps=args.fps, frame_width=args.frame_width
837     )
838     specs.append(
839         {
840             "name": name,
841             "trajectory": traj,
842             "labels": local_label,
843             "fps": args.fps,
844             "frame_width": args.frame_width,
845         }
846     )
847
848 if len(specs) < 2:
849     print(
850         f"[worker] need >=2 videos with trajectories for leave-one-video-out; got
{len(specs)})."
851     )
852     return 2
853 manifest_path = os.path.join(scratch, "manifest.json")
854 with open(manifest_path, "w", encoding="utf-8") as f:
855     json.dump(specs, f, indent=2)
856 src_label = (
857     "real detector" if args.trajectory_source == "detector" else "CPU-mock stub"
858 )
859 print(f"[worker] built manifest: {len(specs)} videos ({src_label} trajectories)")
860
861 out_dir = os.path.join(scratch, "artifacts")
862 cmd = [
863     sys.executable,
864     "-m",
865     "training.gen0.harness",
866     "--manifest",
867     manifest_path,
868     "--out",
869     out_dir,
870     "--version",
871     args.version,
872 ]
873 if args.floor:
874     floor_local = (
875         storage.fetch(
876             args.floor, os.path.join(scratch, "floor.json"), config=storage_cfg
877         )
878         if "://" in args.floor
879         else args.floor
880     )
881     cmd += ["--floor", floor_local]
882 print("[worker] training:", " ".join(cmd))
883 rc = subprocess.run(cmd).returncode
884 if rc != 0:
885     print(f"[worker] gen0 harness failed (rc={rc})")

```

```

886         return rc
887
888     # PUSH the versioned artifact bundle (weights.json + manifest.json) to the bucket.
889     bundle = os.path.join(out_dir, args.version)
890     out_prefix = args.out_uri.rstrip("/")
891     pushed = []
892     for fn in sorted(os.listdir(bundle)):
893         uri = f"{out_prefix}/weights/{args.version}/{fn}"
894         storage.put(os.path.join(bundle, fn), uri, config=storage_cfg)
895         pushed.append(uri)
896     print(f"[worker] train done: pushed {len(pushed)} artifact(s):")
897     for u in pushed:
898         print("    ", u)
899     return 0
900
901
902 def build_parser() -> argparse.ArgumentParser:
903     p = argparse.ArgumentParser(
904         prog="python -m backend.worker",
905         description="Storage-aware worker: pull ? run pipeline ? push.",
906     )
907     p.add_argument(
908         "-v",
909         "--verbose",
910         action="store_true",
911         help="DEBUG logging (the native detector command, frame counts, per-stage
912         detail)",
913     )
914     sub = p.add_subparsers(dest="cmd", required=True)
915
916     pi = sub.add_parser("infer", help="run inference on one video URI and push outputs")
917     pi.add_argument(
918         "--video-uri", required=True, help="local path / local:// / s3:// / gcs://"
919     )
920     pi.add_argument(
921         "--out-uri",
922         required=True,
923         help="output prefix (outputs/<video_id>/? is appended)",
924     )
925     pi.add_argument(
926         "--segmenter", default=None, help="default: indexing.default_segmenter"
927     )
928     pi.add_argument(
929         "--config", default=None, help="config.json path (default: ./config.json)"
930     )
931     pi.add_argument(
932         "--scratch", default=None, help="local scratch dir (default: a temp dir)"
933     )
934     pi.add_argument("--max-frames", type=int, default=None)
935     pi.add_argument(
936         "--done-uri",
937         default=None,
938         help="M6: storage URI for a completion MARKER (video_id + rc) written on a "
939         "successful run, so a remote (Cloud Run) dispatcher's bucket-poll terminates. "
940         "Default-OFF: omit it for the normal local run.",
941     )
942     pi.add_argument(
943         "--set",
944         action="append",
945         default=[],
946         metavar="k=v",
947         help="config override, e.g. --set indexing.skip_ai_handoff=true",
948     )
949     pi.add_argument(
950         "--skip-validation",

```



```

950         action="store_true",
951         help="skip the worker's re-validation: the dispatcher (control-plane) already
ran the "
952         "validators with its resolved config and only dispatches on PASS, so re-
validating here is "
953         "redundant, decodes the video again, and can contradict the gate. Default-OFF:
the direct "
954         "`worker infer` CLI still validates.",
955     )
956     pi.set_defaults(func=cmd_infer)
957
958     pc = sub.add_parser(
959         "compile",
960         help="#521: stitch a highlight reel on this worker (L4 NVENC) from a staged
compile-spec",
961     )
962     pc.add_argument(
963         "--spec-uri",
964         required=True,
965         help="storage URI of the control-plane-staged compile-spec JSON (time-pairs +
stitching "
966         "config + source uri + video_id)",
967     )
968     pc.add_argument(
969         "--out-uri",
970         default="",
971         help="output ROOT (gs://? bucket); the reel is mirrored to
highlights/<video_id>/<name>. "
972         "Also forwarded via --set storage.output_root; either resolves the reel path.",
973     )
974     pc.add_argument(
975         "--done-uri",
976         default=None,
977         help="storage URI for the compile completion MARKER (rc + download_url), so a
remote "
978         "dispatcher's bucket-poll terminates. Default-OFF: omit for a direct-CLI
stitch.",
979     )
980     pc.add_argument(
981         "--config", default=None, help="config.json path (default: ./config.json)"
982     )
983     pc.add_argument(
984         "--scratch", default=None, help="local scratch dir (default: a temp dir)"
985     )
986     pc.add_argument(
987         "--set",
988         action="append",
989         default=[],
990         metavar="k=v",
991         help="config override, e.g. --set storage.output_root=gs://bucket",
992     )
993     pc.set_defaults(func=cmd_compile)
994
995     pt = sub.add_parser(
996         "train",
997         help="Gen-0 train-from-bucket (labels [+videos] -> trajectories -> harness)",
998     )
999     pt.add_argument(
1000         "--corpus-uri",
1001         required=True,
1002         help="prefix containing labels/ (+ videos/ for "
1003         "--trajectory-source detector), e.g. s3://bucket or gcs://bucket",
1004     )
1005     pt.add_argument(
1006         "--out-uri",

```

```

1007         required=True,
1008         help="output prefix (weights/<version>/ is appended)",
1009     )
1010     pt.add_argument("--version", required=True, help="semver for the artifact bundle")
1011     pt.add_argument(
1012         "--trajectory-source",
1013         choices=["synth", "detector"],
1014         default="synth",
1015         help="synth = CPU-mock label-consistent trajectories (default, no GPU); "
1016         "detector = REAL shuttle trajectories via the resolved DetectorRunner "
1017         "(needs videos/ in the corpus; honours detector_impl).",
1018     )
1019     pt.add_argument(
1020         "--segmenter",
1021         default="wasb_hybrid",
1022         help="trajectory hybrid whose detector is used for --trajectory-source",
1023     )
1024     pt.add_argument(
1025         "--floor", default=None, help="floor baseline JSON (local path or storage URI)"
1026     )
1027     pt.add_argument("--config", default=None)
1028     pt.add_argument("--scratch", default=None)
1029     pt.add_argument(
1030         "--fps",
1031         type=float,
1032         default=30.0,
1033         help="synth trajectory fps (detector probes the video)",
1034     )
1035     pt.add_argument(
1036         "--frame-width",
1037         type=float,
1038         default=1280.0,
1039         help="synth trajectory frame width (detector probes the video)",
1040     )
1041     pt.add_argument("--set", action="append", default=[], metavar="k=v")
1042     pt.set_defaults(func=cmd_train)
1043     return p
1044
1045
1046     def _setup_logging(verbose: bool) -> None:
1047         """Configure logging for the worker CLI so the pipeline's existing INFO/ERROR
1048         diagnostics are
1049         VISIBLE ? the worker previously ran with no logging config, so every
1050         segmenter/detector log
1051         (detector choice, 'Running native WASB inference', availability failures) was
1052         silently dropped.
1053         That is exactly why a 0-rally cold-start gave nothing to analyse. Configured here
1054         (the CLI
1055         entry), not at import, so importing the module never reconfigures a caller's
1056         logging."""
1057         # force=True: basicConfig is a no-op if the root already has handlers (e.g. a
1058         library
1059         # configured logging at import), which would leave the level at WARNING and re-
1060         suppress the
1061         # diagnostics. Forcing reconfiguration means the worker reliably controls its own
1062         verbosity.
1063         logging.basicConfig(
1064             level=logging.DEBUG if verbose else logging.INFO,
1065             format="%(asctime)s [%(levelname)s] %(name)s: %(message)s",
1066             datefmt="%H:%M:%S",
1067             force=True,
1068         )

```

```

1063 def main(argv=None) -> int:
1064     args = build_parser().parse_args(argv)
1065     _setup_logging(getattr(args, "verbose", False))
1066     return args.func(args)
1067
1068
1069 if __name__ == "__main__":
1070     sys.exit(main())
1071

```

14. assistant (2026-07-09T06:25:12.861Z)

Now I have the origin/master version with the `\_progress` callback (lines 388?428) and telemetry (lines 359?368). Let me read the three other key files: the native runner (where the callback should be invoked), the telemetry module, and the validator base (for the result structure).

15. user (2026-07-09T06:25:13.957Z)

```

1      """Linux-native WASB shuttle-trajectory runner (no WSL).
2
3      The de-risk goal the ``DetectorRunner`` ABC was built for (see ``base.py`` ?
4      "Optional Linux-native path"; ``docs/DECOUPLED_COMPUTE/04-BUILD-AND-TEST-PLAN.md``
M3.5).
5      It is the sibling of :class:`~backend.pipeline.detectors.wasb_runner.WasbRunner` for a
6      Linux GPU box (GCP / a SaaS worker): it runs the same verified ``wasb_infer.py``
core
7      that the WSL adapter runs, so the trajectory CSV is identical (to a sub-pixel float
tolerance;
8      see the acceptance gate below) ? but WITHOUT the WSL transport:
9
10     * NO ``WslCommandMixin`` (no ``wsl -d <distro> bash -c``), NO ``conda activate`` (the
11     ``wasb`` env's python is invoked by path), NO ``/mnt/c`` path translation
12     (``to_wsl_mnt_path``). Paths are native; the video is already local on the box.
13     * ``device`` is ``cuda`` (default ? the WSL parity path), ``mps``, ``cpu``, or
``auto``
14     (M4: cuda-if-available-else-cpu ? the CPU fallback for a no-GPU box; explicit
``cuda`` stays
15     strict); threaded into ``wasb_infer.py --device`` (and so the Hydra
``runner.device`` override).
16     * weights / repo / python / scratch come from the environment first
17     (``WASB_WEIGHTS`` / ``WASB_REPO_DIR`` / ``WASB_PYTHON`` / ``WASB_SCRATCH``), then
18     ``indexing.wasb.*`` config, then a default ? so a box needs no config edits.
19
20     It emits the identical ``Frame,Visibility,X,Y`` CSV (same ``{stem}__vid12_wasb.csv``
21     name as the WSL runner) and reuses the model-agnostic ``parse_trajectory_csv`` /
22     ``trajectory_to_action_windows`` from ``tracknet_runner``, so the trajectory hybrids and
23     the windowing brain are unchanged. Selected via the existing ``detector_impl`` seam
24     (``trajectory_hybrid._resolve_runner`` ? ``_build_native_runner``);
``detector_impl="native"``.
25
26     Acceptance gate (parity, NOT strict byte-equality): the SAME clip through the WSL
runner
27     (Windows) and this runner (Linux, ``device=cuda``) yields the same trajectory CSV + the
same
28     candidate windows within a small float tolerance (~1e-3 px). cuDNN is non-
deterministic
29     ACROSS GPUs, so a sub-pixel diff is expected on different hardware ? but it is byte-
exact
30     run-to-run on a FIXED GPU. ? Validated 2026-06-20 on an RTX 2070 (via WSL): native
streaming
31     vs the cached WSL-disk CSV matched 1194/1195 frames (the lone diff a ~5e-5 px edge
32     artifact), and two native runs were byte-identical (deterministic). This is a hardware
check
33     (a real GPU + the ``wasb`` env), so the offline, subprocess-mocked suite asserts the
contract,

```

```

34     not the numbers.
35     """
36
37     import logging
38     import os
39     import re
40     import shutil
41     import subprocess
42     from dataclasses import dataclass
43     from typing import Callable, List, Optional, Tuple
44
45     from backend.config.models import IndexingConfig
46     from backend.pipeline.detectors.base import DetectorRunner
47     from backend.pipeline.detectors.device import AUTO, VALID_DEVICES, DeviceContext
48
49     # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical to
WSL/TrackNet.
50     from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
51
52     logger = logging.getLogger(__name__)
53
54     # Path to the inference core we copy into the WASB repo's src/ (so it can import WASB's
55     # detectors/trackers/dataloaders + find configs/), exactly as the WSL adapter does.
56     _INFER_PY = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")
57
58     # Single torch/cuda probe code, used by BOTH healthcheck and the lazy device-resolution
probe
59     # so the "cuda True/False" string they parse can never drift apart.
60     _CUDA_PROBE_CODE = (
61         "import torch; print('torch', torch.__version__, 'cuda', torch.cuda.is_available())"
62     )
63
64     # Decode backends wasb_infer.py understands (#506). Mirrored by its --decode-backend
choices;
65     # healthcheck rejects anything else up front so a typo can't reach argparse as rc=2 mid-
run.
66     _VALID_DECODE_BACKENDS = ("cpu", "nvdec_resident")
67
68     # Probe for the GPU-resident decode dep (#506): a wasb env without PyNvVideoCodec must
fail the
69     # healthcheck, not the paid inference subprocess. (The NVDEC driver libs ? libnvcuid ?
are only
70     # exercised at decoder INIT, so a missing lib still surfaces in the subprocess; this
catches the
71     # common misconfig, a wasb env built without the package.)
72     _PYNVC_PROBE_CODE = "import PyNvVideoCodec"
73
74
75     @dataclass
76     class NativeWasbConfig:
77         """Resolved settings for a Linux-native WASB run.
78
79         Built via :meth:`from_indexing_cfg` with **environment overrides** taking precedence
80         over ``indexing.wasb.*`` config (the box sets env; no config edits needed).
``device``
81         defaults to ``cuda`` (the WSL parity path); ``stream_video`` defaults to True (the
perf win).
82         """
83
84         repo_dir: str = "~/models/WASB-SBDT" # native clone of nttcom/WASB-SBDT
85         python_bin: str = (
86             "python" # the `wasb` conda env python (invoked by path, no activate)
87         )
88         weights_path: str = "" # REQUIRED (env WASB_WEIGHTS or indexing.wasb.weights_path)
89         sport: str = "badminton"

```

```

90     device: str = "cuda" # auto | cuda | mps | cpu (auto = cuda-if-available-else-cpu)
91     scratch_dir: str = "" # frame-extraction scratch (env); "" = next to the output dir
92     timeout_sec: int = 1800 # 30 min; a hung GPU call is killed (0 = no timeout)
93     keep_frames: bool = (
94         False # retain the decoded frame cache after success (debug only)
95     )
96     # NATIVE DEFAULT = STREAMING: skip the cv2 PNG extraction that dominates a long
clip's
97     # wall-clock (measured ~50s on a 20s clip) ? bit-identical to disk (verified on a
real GPU
98     # 2026-06-20). The GPU box reads from a local NVMe so streaming is the right default
here.
99     stream_video: bool = True
100    # GPU-feed tuning (indexing.wasb.batch_size / prefetch_batches). 0 = omit the flag ?
101    # wasb_infer.py's parity defaults (batch 8, no prefetch). The cloud-serving preset
opts in.
102    batch_size: int = 0
103    prefetch_batches: int = 0
104    # GPU-resident NVDEC decode (#506, indexing.wasb.decode_backend /
WASB_DECODE_BACKEND).
105    # "cpu" (default) = omit the flag ? wasb_infer.py's historical cv2 decode paths,
argv
106    # byte-identical. "nvdec_resident" = decode+resize+normalize on-GPU (torch-2.x env
only).
107    decode_backend: str = "cpu"
108    log_every_batches: int = 50 # emission rate of "D/T (pct%)" progress lines from
wasb_infer
109
110    @classmethod
111    def from_indexing_cfg(cls, idx_cfg: "IndexingConfig") -> "NativeWasbConfig":
112        if isinstance(idx_cfg, dict):
113            idx_cfg = IndexingConfig(**idx_cfg)
114            w = idx_cfg.wasb
115            env = os.environ.get
116            # stream_video is tri-state in config: None/absent => native default (stream);
an
117            # explicit True/False still wins (e.g. False for a container cv2 can't seek).
118            sv = w.stream_video
119            return cls(
120                repo_dir=env("WASB_REPO_DIR") or w.repo_dir,
121                python_bin=env("WASB_PYTHON") or w.python_bin,
122                weights_path=env("WASB_WEIGHTS") or w.weights_path or "",
123                sport=w.sport,
124                device=env("WASB_DEVICE") or w.device,
125                scratch_dir=env("WASB_SCRATCH")
126                or env("RALLY_SCRATCH")
127                or env("TMPDIR")
128                or "",
129                timeout_sec=w.timeout_sec,
130                keep_frames=w.keep_frames,
131                stream_video=(True if sv is None else bool(sv)),
132                batch_size=int(w.batch_size or 0),
133                prefetch_batches=int(w.prefetch_batches or 0),
134                decode_backend=env("WASB_DECODE_BACKEND") or w.decode_backend,
135                log_every_batches=int(w.log_every_batches or 50),
136            )
137
138
139    class NativeWasbRunner(DetectorRunner):
140        """WASB runner that executes natively on a Linux GPU box ? no WSL. See module
docstring."""
141
142        def __init__(self, cfg: NativeWasbConfig):
143            self.cfg = cfg
144            # Cached CUDA-availability probe (used to resolve device="auto"). Set by

```

```

healthcheck;
145         # lazily probed by run_predict if healthcheck was skipped. None = not yet
probed.
146         self._cuda_available: Optional[bool] = None
147
148         # --- low-level native exec (the single place tests patch)
-----
149     def _run(
150         self, argv: List[str], cwd: Optional[str] = None
151     ) -> subprocess.CompletedProcess:
152         logger.debug("native exec: %s (cwd=%s)", " ".join(argv), cwd)
153         return subprocess.run(
154             argv,
155             cwd=cwd,
156             capture_output=True,
157             text=True,
158             timeout=(self.cfg.timeout_sec or None),
159         )
160
161     # --- device resolution (M4 DeviceContext: device="auto" ? cuda-if-available-else-
cpu) -----
162     def _probe_cuda(self) -> bool:
163         """Probe ``torch.cuda.is_available()`` in the wasb env's python (a subprocess).
Any failure
164         (missing python / timeout / import error) is treated as 'no CUDA'."""
165         try:
166             res = self._run([self.cfg.python_bin, "-c", _CUDA_PROBE_CODE])
167         except (FileNotFoundError, subprocess.TimeoutExpired):
168             return False
169         return res.returncode == 0 and "cuda True" in (res.stdout or "")
170
171     def _cuda_is_available(self) -> bool:
172         """CUDA availability, cached across healthcheck + run_predict so we probe at
most once."""
173         if self._cuda_available is None:
174             self._cuda_available = self._probe_cuda()
175         return self._cuda_available
176
177     def _effective_device(self) -> str:
178         """The torch device to actually pass to ``wasb_infer.py``. Only ``auto`` needs
the CUDA
179         probe; an explicit cuda/mps/cpu resolves to itself (so ``device='cuda'`` is
unchanged)."""
180         if self.cfg.device == AUTO:
181             return DeviceContext(AUTO, self._cuda_is_available()).effective
182         return self.cfg.device
183
184     def _repo_src(self) -> str:
185         return os.path.join(os.path.expanduser(self.cfg.repo_dir), "src")
186
187     def _copy_infer_script(self) -> bool:
188         """Copy wasb_infer.py into the WASB repo src/ so it imports WASB modules + finds
configs/."""
189         repo_src = self._repo_src()
190         try:
191             os.makedirs(repo_src, exist_ok=True)
192             shutil.copy(_INFER_PY, os.path.join(repo_src, "wasb_infer.py"))
193             return True
194         except OSError as e:
195             logger.error("Failed to copy wasb_infer.py into %s: %s", repo_src, e)
196             return False
197
198     def _rm_rf(self, path: Optional[str]) -> None:
199         """Best-effort recursive delete of a native path (post-success cleanup; never
raises)."""

```

```

200         if path:
201             shutil.rmtree(path, ignore_errors=True)
202
203     def healthcheck(self) -> Tuple[bool, str]:
204         """Verify weights + repo present and the `wasb` python imports torch (and sees a
GPU on cuda)."""
205         if not self.cfg.weights_path:
206             return False, (
207                 "WASB weights path is empty ? set WASB_WEIGHTS (or
indexing.wasb.weights_path)."
208             )
209         weights = os.path.expanduser(self.cfg.weights_path)
210         if not os.path.exists(weights):
211             return False, f"WASB weights not found at {weights}"
212         repo_src = self._repo_src()
213         if not os.path.isdir(repo_src):
214             return False, (
215                 f"WASB-SBDT repo src/ not found at {repo_src} ? clone nttcom/WASB-SBDT "
216                 f"({set WASB_REPO_DIR / indexing.wasb.repo_dir})."
217             )
218         # Fail fast on a typo'd device (the device-policy seam) ? a clear error here
beats a cryptic
219         # argparse rc=2 from wasb_infer.py at run time.
220         if self.cfg.device not in VALID_DEVICES:
221             return False, (
222                 f"unknown device {self.cfg.device!r} ? valid: {'',
'.join(VALID_DEVICES)})."
223             )
224         if self.cfg.decode_backend not in _VALID_DECODE_BACKENDS:
225             return False, (
226                 f"unknown decode_backend {self.cfg.decode_backend!r} ? valid: "
227                 f"{'', '.join(_VALID_DECODE_BACKENDS)})."
228             )
229         try:
230             res = self._run([self.cfg.python_bin, "-c", _CUDA_PROBE_CODE])
231         except FileNotFoundError:
232             return (
233                 False,
234                 f"python binary not found: {self.cfg.python_bin!r} (set WASB_PYTHON).",
235             )
236         except subprocess.TimeoutExpired:
237             return False, "native torch healthcheck timed out."
238         out = (res.stdout or "").strip()
239         if res.returncode != 0:
240             return (
241                 False,
242                 f"`{self.cfg.python_bin}` torch import failed: {(res.stderr or
out).strip()}",
243             )
244         # M4: resolve the device. Cache the probe so run_predict reuses it (no second
subprocess).
245         self._cuda_available = "cuda True" in out
246         ctx = DeviceContext(self.cfg.device, self._cuda_available)
247         if ctx.cuda_required_but_missing:
248             # Unchanged pre-M4 behaviour for an EXPLICIT cuda request: hard-fail, never
a silent
249             # slow-CPU downgrade. Set indexing.wasb.device=auto for a CPU fallback.
250             return False, f"device=cuda but torch.cuda.is_available() is False ({out})"
251         if ctx.fell_back:
252             logger.warning(
253                 "device=auto: no CUDA GPU detected (%s) ? FALLING BACK TO CPU; WASB "
254                 "inference will be slow. Set indexing.wasb.device=cuda to require a
GPU.",
255                 out,
256             )

```

```

257         # NVDEC decode is CUDA-only: an auto?cpu fallback (or explicit cpu/mps, which
the
258         # config validator already rejects) cannot run the GPU-resident path ? fail loud
259         # instead of letting wasb_infer.py die mid-run on a paid box.
260         if self.cfg.decode_backend == "nvdec_resident" and ctx.effective != "cuda":
261             return False, (
262                 f"decode_backend=nvdec_resident requires CUDA, but the effective device
is "
263                 f"{ctx.effective!r} ({out}) ? NVDEC decode runs on the GPU."
264             )
265         if self.cfg.decode_backend == "nvdec_resident":
266             try:
267                 nv = self._run([self.cfg.python_bin, "-c", _PYNVC_PROBE_CODE])
268             except (FileNotFoundError, subprocess.TimeoutExpired):
269                 return False, "PyNvVideoCodec probe failed to run (python
missing/timeout)."
270             if nv.returncode != 0:
271                 return False, (
272                     f"decode_backend=nvdec_resident but `{_PYNVC_PROBE_CODE}` failed in
"
273                     f"{self.cfg.python_bin!r}: {(nv.stderr or nv.stdout or
''.strip())[-300:]}"
274                     f"? the wasb env needs PyNvVideoCodec (deploy/cloudrun/Dockerfile /
"
275                     f"gpu_setup.sh stage it)."
276                 )
277             return True, out
278
279     def run_predict(
280         self,
281         video_win_path: str,
282         output_win_dir: str,
283         video_id: Optional[str] = None,
284         progress: Optional[Callable[[str], None]] = None,
285     ) -> Optional[str]:
286         """Run WASB natively on a video. Returns the path to the trajectory CSV, or
None.
287
288         Output artifacts are namespaced by ``video_id`` (the file MD5) exactly as the
WSL
289         runner does, and the CSV is named ``{stem}__{vid12}_wasb.csv`` ? byte-for-byte
the
290         same downstream contract, so the parity comparison is apples-to-apples.
291         """
292         output_dir = os.path.abspath(output_win_dir)
293         os.makedirs(output_dir, exist_ok=True)
294         # Normalize for a cross-platform basename (tests pass Windows paths on Linux).
295         norm = str(video_win_path).replace("\\", "/")
296         stem, _ext = os.path.splitext(os.path.basename(norm))
297         # Content-derived key: same filename + different bytes -> different key (no
cache reuse).
298         key = f"{stem}__{video_id[:12]}" if video_id else (stem or "wasb")
299         out_csv = os.path.join(output_dir, f"{key}_wasb.csv")
300
301         if not self._copy_infer_script():
302             logger.error("Could not stage wasb_infer.py into the WASB repo src/.")
303             return None
304
305         weights = os.path.expanduser(self.cfg.weights_path)
306         argv = [
307             self.cfg.python_bin,
308             "wasb_infer.py",
309             "--video",
310             os.path.abspath(str(video_win_path)),
311         ]

```



```

312 frames_dir: Optional[str] = None
313 if self.cfg.decode_backend == "nvdec_resident":
314     # GPU-resident path (#506): wasb_infer.py decodes the video on the GPU
315     # (PyNvVideoCodec NVDEC) ? no PNG extraction and no cv2 streaming, so
316     # --stream-video nor --frames_out_dir applies and there is no frame cache to
317     # clean.
318     argv += ["--decode-backend", self.cfg.decode_backend]
319 elif self.cfg.stream_video:
320     # Fast path: decode in-process, no PNG extraction. Output is bit-identical
321     # to disk.
322     argv.append("--stream-video")
323 else:
324     scratch = self.cfg.scratch_dir or output_dir
325     frames_dir = os.path.join(scratch, f"{key}_frames")
326     argv += ["--frames_out_dir", frames_dir]
327 # M4: resolve device="auto" to the effective device (cuda-if-available-else-
328 # cpu); an
329 # explicit cuda/mps/cpu passes through unchanged. Never sends "auto" to
330 wasb_infer.py.
331 device = self._effective_device()
332 if self.cfg.decode_backend == "nvdec_resident" and device != "cuda":
333     # Mirrors the healthcheck gate for callers that skipped it (e.g. auto?cpu on
334     # a
335     # GPU-less box): NVDEC decode cannot run without CUDA ? fail before the
336     # subprocess.
337     logger.error(
338         "decode_backend=nvdec_resident requires CUDA, but the effective device
339         is "
340         "%r ? aborting before the WASB subprocess.",
341         device,
342     )
343     return None
344 argv += [
345     "--weights",
346     weights,
347     "--sport",
348     self.cfg.sport,
349     "--device",
350     device,
351     "--out",
352     out_csv,
353 ]
354 # GPU-feed tuning: pass only explicit non-zero values so the default argv (and
355 # its
356 # byte-parity guarantee) is unchanged when the knobs are unset.
357 if self.cfg.batch_size:
358     argv += ["--batch-size", str(self.cfg.batch_size)]
359 if self.cfg.prefetch_batches and self.cfg.decode_backend != "nvdec_resident":
360     # Prefetch overlaps CPU-side batch assembly with GPU compute ? meaningless
361     # for
362     # the GPU-resident decoder (there is no CPU assembly), whose wasb_infer
363     # branch
364     # ignores the flag; keep the argv honest instead of passing a dead knob.
365     argv += ["--prefetch-batches", str(self.cfg.prefetch_batches)]
366 if self.cfg.log_every_batches and self.cfg.log_every_batches > 0:
367     argv += ["--log-every-batches", str(self.cfg.log_every_batches)]
368
369 logger.info(
370     "Running native WASB inference (device=%s%s) on %s ...",
371     device,
372     " [auto?fallback]" if self.cfg.device == AUTO and device != "cuda" else "",
373     os.path.basename(norm),
374 )

```

```

365
366 res: Optional[subprocess.CompletedProcess] = None
367 try:
368     if progress is None:
369         # Fast path: capture at end (existing behaviour, byte-identical for
tests).
370         res = self._run(argv, cwd=self._repo_src())
371     else:
372         # Live progress path (#534): stream stdout so we can parse the periodic
373         # "done/total (pct%) ..." lines already emitted by wasb_infer.py and
call
374         # the callback. This gives the SPA a real heartbeat during the long GPU
pass
375         # instead of a spinner. Still respects timeout.
376         proc = subprocess.Popen(
377             argv,
378             cwd=self._repo_src(),
379             stdout=subprocess.PIPE,
380             stderr=subprocess.PIPE,
381             text=True,
382         )
383         assert proc.stdout is not None
384         try:
385             for line in proc.stdout:
386                 line = line.rstrip()
387                 if line:
388                     # Parse existing wasb_infer progress lines for detector
heartbeat.
389                     # Example: "1234/5000 (24%) | 2.3 win/s | elapsed 45s | ETA
120s"
390                     m = re.search(r"(\d+)/(\d+)\s+\((\d+)\%\)\"", line)
391                     if m:
392                         done, total, pct = m.groups()
393                         try:
394                             progress(f"detector {pct}% ({done}/{total})")
395                         except Exception: # noqa: BLE001
396                             pass # progress must never kill the run
397                         # Also log for the container/runlog.
398                         logger.info(line)
399                     # Wait for completion (communicate for stderr/returncode).
400                     _, stderr = proc.communicate()
401                     res = subprocess.CompletedProcess(
402                         args=argv, returncode=proc.returncode, stdout="", stderr=stderr
or ""
403                 )
404             finally:
405                 if proc.poll() is None:
406                     proc.kill()
407         except FileNotFoundError:
408             logger.error(
409                 "python binary not found: %r (set WASB_PYTHON).", self.cfg.python_bin
410             )
411             return None
412         except subprocess.TimeoutExpired:
413             logger.error(
414                 "WASB native inference timed out after %ss.", self.cfg.timeout_sec
415             )
416             return None
417
418 if res is None or res.returncode != 0:
419     logger.error(
420         "WASB native inference failed:\n%s",
421         (getattr(res, "stderr", None) or getattr(res, "stdout", None) or
"")[-2000:],
422     )

```

```

423         return None
424     if not os.path.exists(out_csv):
425         logger.error(
426             "WASB native finished but expected CSV not found at %s", out_csv
427         )
428         return None
429     logger.info("WASB trajectory CSV: %s", out_csv)
430
431     # Success ? the CSV is the durable output; the decoded frame cache is a pure
intermediate.
432     # On any earlier failure we return above WITHOUT cleaning, so the cache survives
a resume.
433     # (Streaming never materializes frames, so there is nothing to drop on that
path.)
434     if not self.cfg.keep_frames and frames_dir:
435         logger.info(
436             "Cache hygiene: removing frame cache for %s "
437             "(set indexing.wasb.keep_frames=true to retain).",
438             key,
439         )
440         self._rm_rf(frames_dir)
441     return out_csv
442
443     # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
444     parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
445     trajectory_to_action_windows = staticmethod(
446         TrackNetRunner.trajectory_to_action_windows
447     )
448

```

16. assistant (2026-07-09T06:25:48.732Z)

The native runner's streaming path (lines 366-406) reads `**only `proc.stdout`**` and parses ``(\d+)/(\d+)\s+((\d+)\%)\``. The critical question for GAP 3: how does ``wasb_infer.py`` actually emit its progress lines ? to stdout via ``print``, or to stderr via ``logging``? And does the ``nvdec_resident`` branch (the live serving default) emit them at all? Let me read ``wasb_infer.py``, the segmenter plumbing, the telemetry module, and the validator base.

17. user (2026-07-09T06:25:50.274Z)

```

1     """Structured run telemetry ? a compact, rotating JSONL "black box" for long/flaky runs.
2
3     A long GPU/CV run (the fusion-golden extraction, a WSL/WASB pass, an ffmpeg stitch) can
die with
4     **no Python traceback** ? the OS kills it (a console-control event, an OOM, a thermal
event). When
5     that happens the only forensic evidence is whatever was last flushed to disk. This
module keeps a
6     tiny (~10 KB) rotating runlog of *execution velocity + machine health* so a post-mortem
is a
7     bullseye instead of a guess: "it was at 26 fps, GPU 85 °C throttled to 1125/2100 MHz,
4.8 GB VRAM
8     free, CPU 60 %, RAM 21/32 GB ? then vanished."
9
10    It is the **observability** companion to the `ExecutionPolicy` *control* layer
11    (`docs/EXECUTION_POLICY.md`): the policy decides wait / abort / resume; this records the
signals
12    that explain *why* a run slowed or died, and lets us compare execution velocity across
quality
13    iterations.
14
15    Design:
16    - **Caller-driven snapshots** (no background thread): call `snapshot(done, total)`
from an

```

```

17     existing progress callback; fps is derived from the delta since the previous snapshot.
18     ``event(kind, msg)`` records discrete moments (start, video-done, error) WITH a health
sample, so
19     the last line before a kill captures the machine state at death.
20     - **Capped + rotating**: a fixed ``meta`` header line (machine / GPU / start) is always
kept; recent
21     snapshot/event lines form an on-disk ring trimmed to ``max_bytes`` (~10 KB). The file
is rewritten
22     atomically (tmp + ``os.replace``) on each write, so a hard kill never leaves it half-
written.
23     - **Graceful degradation**: GPU stats come from ``nvidia-smi`` (absent -> ``None``);
system stats
24     from ``psutil`` (not installed -> ``None``). Collection never raises into the caller ?
a telemetry
25     hiccup must not break the run it observes.
26     - **Deterministic / testable**: the clock and the two collectors are injectable, so fps,
rotation,
27     and parsing are unit-tested with no real GPU, no psutil, and no wall-clock.
28
29     All stdlib + optional ``psutil`` (BSD-3, commercial-clean). No network, no training-data
path.
30     """
31
32     from __future__ import annotations
33
34     import json
35     import os
36     import shutil
37     import subprocess
38     import time
39     from typing import Any, Callable, Dict, List, Optional
40
41     GpuFn = Callable[[], Optional[Dict[str, Any]]]
42     SysFn = Callable[[], Optional[Dict[str, Any]]]
43
44     # nvidia-smi fields pulled in one query, in this order (see gpu_stats parsing).
45     _GPU_QUERY =
"temperature.gpu,utilization.gpu,power.draw,clocks.sm,clocks.max.sm,memory.used,memory.free"
46     _THROTTLE_CLOCK_FRAC = 0.8 # SM clock below this fraction of max => flagged throttled
47     _THROTTLE_TEMP_C = (
48         84.0 # ... or temperature at/above this (laptop Max-Q throttle territory)
49     )
50
51
52     def _num(s: str) -> Optional[float]:
53         try:
54             return float(s)
55         except (TypeError, ValueError):
56             return None
57
58
59     def gpu_stats(timeout: float = 4.0) -> Optional[Dict[str, Any]]:
60         """One ``nvidia-smi`` sample as a dict, or ``None`` if nvidia-smi is unavailable /
errors.
61
62         Returns util %, temperature, power (W), current + max SM clock (MHz), VRAM used /
free (MB), and a
63         derived ``throttled`` flag (SM clock well below max, or hot). Pure read; never
raises.
64         """
65         exe = shutil.which("nvidia-smi")
66         if not exe:
67             return None
68         try:
69             out = subprocess.run(

```

```

70         [exe, f"--query-gpu={_GPU_QUERY}", "--format=csv,noheader,nounits"],
71         capture_output=True,
72         text=True,
73         timeout=timeout,
74         check=True,
75     ).stdout.strip()
76 except (subprocess.SubprocessError, OSError):
77     return None
78 row = out.splitlines()[0] if out else ""
79 parts = [p.strip() for p in row.split(",")]
80 if len(parts) < 7:
81     return None
82 vals = [_num(p) for p in parts[:7]]
83 temp, util, power, sm, sm_max, used, free = vals
84 throttled = bool(temp is not None and temp >= _THROTTLE_TEMP_C)
85 if sm is not None and sm_max:
86     throttled = throttled or sm < _THROTTLE_CLOCK_FRAC * sm_max
87 return {
88     "util_pct": util,
89     "temp_c": temp,
90     "power_w": power,
91     "sm_mhz": sm,
92     "sm_max_mhz": sm_max,
93     "vram_used_mb": used,
94     "vram_free_mb": free,
95     "throttled": throttled,
96 }
97
98
99 def system_stats() -> Optional[Dict[str, Any]]:
100     """CPU % + RAM via ``psutil``, or ``None`` if psutil isn't installed. Never
101     raises."""
102     try:
103         import psutil # optional dep (BSD-3); only used when present
104     except ImportError:
105         return None
106     try:
107         vm = psutil.virtual_memory()
108         return {
109             "cpu_pct": psutil.cpu_percent(interval=None),
110             "ram_used_gb": round(vm.used / 1e9, 2),
111             "ram_total_gb": round(vm.total / 1e9, 2),
112             "ram_pct": vm.percent,
113         }
114     except Exception: # noqa: BLE001 - psutil edge cases must not break the run
115         return None
116
117 def gpu_name() -> Optional[str]:
118     """The GPU model string (for the runlog header), or ``None`` without nvidia-smi."""
119     exe = shutil.which("nvidia-smi")
120     if not exe:
121         return None
122     try:
123         out = subprocess.run(
124             [exe, "--query-gpu=name", "--format=csv,noheader"],
125             capture_output=True,
126             text=True,
127             timeout=4.0,
128             check=True,
129         ).stdout
130     except (subprocess.SubprocessError, OSError):
131         return None
132     return out.splitlines()[0].strip() if out else None
133

```

```

134
135     class RunTelemetry:
136         """A small rotating JSONL runlog of execution velocity + machine health.
137
138         Call ``snapshot(done, total, phase=...)`` from a progress callback and ``event(kind,
msg)`` at
139         milestones. The file holds a ``meta`` header line plus the most recent snapshot /
event lines
140         that fit within ``max_bytes`` (~10 KB), rewritten atomically on every write.
141         """
142
143         def __init__(
144             self,
145             path: str,
146             label: str = "run",
147             max_bytes: int = 10_000,
148             gpu_fn: GpuFn = gpu_stats,
149             sys_fn: SysFn = system_stats,
150             clock: Callable[[], float] = time.time,
151         ) -> None:
152             self.path = path
153             self.label = label
154             self.max_bytes = max_bytes
155             self._gpu_fn = gpu_fn
156             self._sys_fn = sys_fn
157             self._clock = clock
158             self._t0 = clock()
159             self._last_done: Optional[int] = None
160             self._last_t: Optional[float] = None
161             self._lines: List[
162                 str
163             ] = [] # recent body lines (on-disk ring, trimmed to max_bytes)
164             self._meta = json.dumps(
165                 {
166                     "meta": {
167                         "label": label,
168                         "started": self._iso(self._t0),
169                         "pid": os.getpid(),
170                         "gpu_name": gpu_name(),
171                         "cpu_count": os.cpu_count(),
172                     }
173                 },
174                 separators=(",", ":"),
175             )
176             self._write()
177
178         # -- public API -----
179         def snapshot(
180             self,
181             done: Optional[int] = None,
182             total: Optional[int] = None,
183             phase: Optional[str] = None,
184             **extra: Any,
185         ) -> Dict[str, Any]:
186             """Record a velocity + health sample; ``fps`` is derived from the previous
snapshot."""
187             now = self._clock()
188             rec: Dict[str, Any] = {
189                 "t": self._iso(now),
190                 "elapsed_s": round(now - self._t0, 1),
191             }
192             if phase is not None:
193                 rec["phase"] = phase
194             if done is not None:
195                 rec["done"] = done

```

```

196         if total:
197             rec["total"] = total
198             rec["pct"] = round(100.0 * done / total, 1)
199         if (
200             self._last_done is not None
201             and self._last_t is not None
202             and now > self._last_t
203         ):
204             rec["fps"] = round((done - self._last_done) / (now - self._last_t), 1)
205             self._last_done, self._last_t = done, now
206         rec.update(extra)
207         self._attach_health(rec)
208         self._append(rec)
209         return rec
210
211     def event(self, kind: str, msg: str = "", **extra: Any) -> Dict[str, Any]:
212         """Record a milestone (start, video_done, error, restart) WITH a health
sample."""
213         now = self._clock()
214         rec: Dict[str, Any] = {
215             "t": self._iso(now),
216             "elapsed_s": round(now - self._t0, 1),
217             "event": kind,
218         }
219         if msg:
220             rec["msg"] = msg
221         rec.update(extra)
222         self._attach_health(rec)
223         self._append(rec)
224         return rec
225
226     # -- internals -----
227     def _attach_health(self, rec: Dict[str, Any]) -> None:
228         gpu = self._safe(self._gpu_fn)
229         sys_ = self._safe(self._sys_fn)
230         if gpu is not None:
231             rec["gpu"] = gpu
232         if sys_ is not None:
233             rec["sys"] = sys_
234
235     def _append(self, rec: Dict[str, Any]) -> None:
236         self._lines.append(json.dumps(rec, separators=(",", ":")))
237         budget = self.max_bytes - len(self._meta) - 1
238         while len(self._lines) > 1 and sum(len(s) + 1 for s in self._lines) > budget:
239             self._lines.pop(0) # drop oldest body line; the meta header is always kept
240         self._write()
241
242     def _write(self) -> None:
243         try:
244             tmp = self.path + ".tmp"
245             with open(tmp, "w", encoding="utf-8") as f:
246                 f.write(self._meta + "\n")
247                 for s in self._lines:
248                     f.write(s + "\n")
249             os.replace(tmp, self.path)
250         except OSError:
251             pass # telemetry must never break the run it observes
252
253     @staticmethod
254     def _safe(fn: GpuFn) -> Optional[Dict[str, Any]]:
255         try:
256             return fn()
257         except Exception: # noqa: BLE001 - a collector hiccup must not break the run
258             return None
259

```

```

260     @staticmethod
261     def _iso(t: float) -> str:
262         return time.strftime("%Y-%m-%dT%H:%M:%S", time.localtime(t))
263

```

18. user (2026-07-09T06:25:50.631Z)

```

1     import hashlib
2     import json
3     import logging
4     from abc import ABC, abstractmethod
5     from typing import Any, List
6
7     from pydantic import BaseModel
8
9     from backend.pipeline.validators.frame_sampler import (
10         FRAME_META_KEY,
11         FRAME_STATES_KEY,
12         SampleMeta,
13         run_frame_sampling,
14     )
15
16     logger = logging.getLogger(__name__)
17
18     # Bumped when validator LOGIC (not config) changes the verdict for identical
19     bytes+config ? e.g. a
20     # validator is rewritten, or the sampling arithmetic changes. It rides inside
21     config_fingerprint, so
22     # a bump invalidates every cached validation verdict (see backend/orchestration.py's
23     validation cache
24     # + backend/infrastructure/database.py v10). The validator SET and each validator's
25     config-derived
26     # thresholds are already folded into the fingerprint; this covers changes the config
27     can't express.
28     _VALIDATION_FINGERPRINT_SCHEMA = 1
29
30
31     class ValidationResult(BaseModel):
32         validator_name: str
33         passed: bool
34         message: str
35         details: dict[str, Any] = {}
36
37
38     class VideoValidator(ABC):
39         @property
40         @abstractmethod
41         def name(self) -> str:
42             """Unique identifier for this validator."""
43             pass
44
45         @property
46         @abstractmethod
47         def description(self) -> str:
48             """Human-readable explanation of what this validator checks."""
49             pass
50
51         @abstractmethod
52         def validate(
53             self, video_path: str, config: Any, context: dict[str, Any]
54         ) -> ValidationResult:
55             """
56             Executes validation.
57             - video_path: Local path to the MP4 file.
58             - config: AppConfig (typed) ? callers pass the typed model; attribute access

```



```

54         preferred over .get() dict chains.
55     - context: Dynamic context dict shared across validators (stores shared values
56       like ffprobe metadata).
57     """
58     pass
59
60
61 class FrameSamplingValidator(VideoValidator):
62     """A validator whose work is sampling decoded frames.
63
64     The three frame-decoding validators (scene_cut, blur, relevance) used to each
65     open their own ``cv2.VideoCapture`` and seek to their sample frames ? ~125
66     random ``POS_FRAMES`` seeks across three decode opens, which cost minutes on
67     long-GOP 4K clips (see ``frame_sampler``). Subclasses instead declare *which*
68     frame indices they want (:meth:`sample_indices`) and reduce those frames
69     online (:meth:`_init_state` / :meth:`_consume`), so the registry can decode
70     the whole set **once** in a single shared forward pass.
71
72     :meth:`validate` uses the shared pass' result when the registry ran it; when
73     a subclass is invoked standalone (direct ``.validate`` call, no shared
74     context) it decodes its own indices in one pass via the same engine. Either
75     way the frames ? and therefore the pass/fail decision ? are identical to the
76     old per-validator seeks.
77     """
78
79     @abstractmethod
80     def sample_indices(self, total_frames: int, fps: float) -> List[int]:
81         """Absolute frame indices this validator samples, in ascending order.
82
83         Pure and deterministic in ``(total_frames, fps)`` ? the shared pass unions
84         these across validators to drive one decode, and the regression suite pins
85         them to the legacy per-validator seek arithmetic.
86         """
87         ...
88
89     @abstractmethod
90     def _init_state(self, config: Any) -> Any:
91         """Fresh accumulator for one run (may pre-resolve config-derived params)."""
92         ...
93
94     @abstractmethod
95     def _consume(self, state: Any, frame_idx: int, frame: Any) -> None:
96         """Fold one sampled frame (BGR ndarray) into ``state``. Called in index
97         order."""
98         ...
99
100    @abstractmethod
101    def _finalize(
102        self, state: Any, meta: SampleMeta, config: Any, context: dict[str, Any]
103    ) -> ValidationResult:
104        """Render the ValidationResult from the accumulated ``state`` + decode meta."""
105        ...
106
107    def validate(
108        self, video_path: str, config: Any, context: dict[str, Any]
109    ) -> ValidationResult:
110        states = context.get(FRAME_STATES_KEY)
111        if states is not None and self.name in states:
112            # Shared single pass already decoded and reduced our frames.
113            meta = context[FRAME_META_KEY]
114            state = states[self.name]
115        else:
116            # Standalone: decode just this validator's indices in one pass.
117            meta, own_states = run_frame_sampling(video_path, [self], config)
118            state = own_states[self.name]

```

```

118         return self._finalize(state, meta, config, context)
119
120
121     class ValidationRegistry:
122         def __init__(self):
123             self._validators: List[VideoValidator] = []
124
125         def register(self, validator: VideoValidator):
126             """Registers a new validator."""
127             self._validators.append(validator)
128
129         def run_all_with_context(self, video_path: str, config: Any):
130             """Runs all validators and returns ``(results, context)``.
131
132             The per-run ``context`` accumulates shared values (notably
133             ``context['ffprobe_meta']`` produced by FormatValidator). ``run_all``
134             discards it; callers that want media metadata for an observability report
135             use this variant instead of re-probing.
136             """
137             results = []
138             context: dict[
139                 str, Any
140             ] = {} # Shared context (e.g. for caching ffprobe outputs)
141             self._prepare_frame_samples(video_path, config, context)
142             for val in self._validators:
143                 try:
144                     res = val.validate(video_path, config, context)
145                     results.append(res)
146                 except Exception as e:
147                     results.append(
148                         ValidationResult(
149                             validator_name=val.name,
150                             passed=False,
151                             message=f"Validator threw an unexpected error: {str(e)}",
152                         )
153                     )
154             return results, context
155
156         def _prepare_frame_samples(
157             self, video_path: str, config: Any, context: dict[str, Any]
158         ) -> None:
159             """Decode every frame-sampling validator's frames in ONE shared pass.
160
161             Populates ``context`` with the shared decode meta + per-validator
162             accumulator state; each :class:`FrameSamplingValidator` then finalizes
163             from that instead of opening its own capture. Best-effort: any failure
164             here is swallowed so validators transparently fall back to a standalone
165             pass (the run must never crash on a decode hiccup)."""
166             samplers = [
167                 v for v in self._validators if isinstance(v, FrameSamplingValidator)
168             ]
169             if not samplers:
170                 return
171             try:
172                 meta, states = run_frame_sampling(video_path, samplers, config)
173             except Exception:
174                 logger.exception(
175                     "Shared frame-sampling pass failed for %s; validators will fall "
176                     "back to standalone decode.",
177                     video_path,
178                 )
179             return
180             context[FRAME_META_KEY] = meta
181             context[FRAME_STATES_KEY] = states
182

```

```

183     def run_all(self, video_path: str, config: Any) -> List[ValidationResult]:
184         """Runs all registered validators in sequence, accumulating results."""
185         results, _ = self.run_all_with_context(video_path, config)
186         return results
187
188     def config_fingerprint(self, config: Any) -> str:
189         """Stable hash of EVERYTHING that determines a validation verdict for fixed
video bytes.
190
191         Powers the video_id-keyed validation-result cache (backend/orchestration.py): a
reupload of
192         identical content reuses its stored verdict ONLY while this fingerprint still
matches, so any
193         change that could flip a decision correctly forces a re-validation (never serve
a stale
194         PASS/FAIL). What goes in ? and why each matters:
195
196         * the whole ``validation`` config block ? every threshold the validators read
197         (min_blur_threshold / min_fps / min_resolution / max_scene_cuts_per_minute /
pts gap /
198         relevance court-color bounds). Hashing the full model (not a hand-picked
subset) means a
199         NEW validation knob is covered automatically.
200         * ``sport`` ? the relevance validator keys its court-color profile off
``config.sport`` (a
201         top-level field, NOT under ``validation``), so switching sports changes the
verdict.
202         * the sorted validator NAMES ? adding/removing a validator changes which checks
run.
203         * a schema version (:data:``_VALIDATION_FINGERPRINT_SCHEMA``) ? a manual escape
hatch for
204         validator LOGIC changes the config can't express.
205
206         Accepts a typed ``AppConfig`` (attribute access) or a plain dict;
``default=str`` keeps the
207         dump total on any exotic value. Deterministic: ``sort_keys`` fixes ordering."""
208         validation: Any = getattr(config, "validation", None)
209         if validation is None and isinstance(config, dict):
210             validation = config.get("validation")
211         validation_repr = (
212             validation.model_dump(mode="json")
213             if hasattr(validation, "model_dump")
214             else validation
215         )
216         sport: Any = getattr(config, "sport", None)
217         if sport is None and isinstance(config, dict):
218             sport = config.get("sport")
219         payload = {
220             "schema": _VALIDATION_FINGERPRINT_SCHEMA,
221             "validators": sorted(v.name for v in self._validators),
222             "sport": sport,
223             "validation": validation_repr,
224         }
225         blob = json.dumps(
226             payload, sort_keys=True, default=str, separators=(",", ":")
227         )
228         return hashlib.sha256(blob.encode("utf-8")).hexdigest()
229
230
231     # Shared global registry
232     registry = ValidationRegistry()
233

```

`RunTelemetry.snapshot(done, total, phase=...)` matches the `\_progress` call site.
`ValidationResult` has `validator\_name/passed/message/details` ? **no top-level `score`** ; the
"score" the fix wants must come out of `details`. Now the two decisive questions for GAP 3: (a)
does `process\_video` actually forward `progress` down to `run\_predict`, and (b) how/where does
`wasb\_infer.py` emit the `done/total (pct%)` lines ? stdout or stderr, and does the nvdec branch
emit them? Let me search.

20. user (2026-07-09T06:26:04.827Z)

```
180:         progress: Optional[Callable[[str], None]] = None,
300:         csv_path = runner.run_predict(video_path, out_dir, video_id=video_id,
progress=progress)
```

21. user (2026-07-09T06:26:06.542Z)

```
22:     ``<cache>/det_raw.jsonl`` (flushed+fsync'd every ``--flush-every`` windows) and
24:     most ``--flush-every`` windows of GPU work instead of the whole pass.
46:         --out /path/traj.csv [--sport badminton] [--limit N] [--cache-dir DIR] [--flush-every
1000] [--fresh]
49: import argparse
168:         f.flush()
277:             break # truncated trailing line (crash mid-flush)
289:         """"Append window records as JSONL and durably flush (bounded loss on crash).""""
292:         fh.flush()
303:         f.flush()
524:         log_every_batches: int = 50,
526:         flush_every: int = 1000,
624:         logger.info(
629:             logger.info("cache incompatible with this run -> starting fresh")
678:         log_every = max(1, log_every_batches)
679:         flush_n = max(1, flush_every)
681:         logger.info(
729:             if len(pending) >= flush_n or done >= windows_total:
734:             if (bi + 1) % log_every == 0 or done >= windows_total:
739:                 logger.info(
741:                     f"({100 * done // max(1, windows_total)})% | {rate:.1f} win/s |
"
742:                     f"elapsed {el:.0f}s | ETA {eta:.0f}s"
750:         logger.info(f"detection done in {time.time() - t0:.0f}s; running tracker...")
752:         logger.info("detector cache already complete -> tracker-only re-run")
783:         logger.info(f"wrote {len(rows)} rows -> {out_csv}")
784:         logger.info(f"visible-shuttle frames: {visible}/{len(rows)}")
806:         logger.info(f"reusing {existing} already-extracted frames -> {frames_dir}")
816:         logger.info(f"extracted {i} frames -> {frames_dir}")
882:         log_every_batches: int = 50,
884:         flush_every: int = 1000,
911:             log_every_batches=log_every_batches,
913:             flush_every=flush_every,
1069:         log_every_batches: int = 50,
1071:         flush_every: int = 1000,
1112:         logger.info(
1136:             logger.info(
1141:                 logger.info("cache incompatible with this run -> starting fresh")
1175:                 log_every = max(1, log_every_batches)
1176:                 flush_n = max(1, flush_every)
1180:                 logger.info(
1220:                     if len(pending) >= flush_n or done >= windows_total:
1225:                     if n_batches % log_every == 0 or done >= windows_total:
1230:                         logger.info(
1232:                             f"({100 * done // max(1, windows_total)})% | {rate:.1f}
win/s | "
1233:                             f"elapsed {el:.0f}s | ETA {eta:.0f}s"
1235:                     if pending: # safety net; the done>=windows_total branch already flushed
1241:                     logger.info(f"detection done in {time.time() - t0:.0f}s; running tracker...")
1243:                     logger.info("detector cache already complete -> tracker-only re-run")
```

```

1250: def main():
1251:     ap = argparse.ArgumentParser()
1297:     "--log-every-batches", type=int, default=50, help="progress log cadence"
1305:     "--flush-every",
1332:         log_every_batches=args.log_every_batches,
1334:         flush_every=args.flush_every,
1349:         log_every_batches=args.log_every_batches,
1351:         flush_every=args.flush_every,
1373:     log_every_batches=args.log_every_batches,
1375:     flush_every=args.flush_every,

```

22. assistant (2026-07-09T06:26:20.781Z)

****Root cause of GAP 3 found.**** ``wasb_infer.py`` emits its progress lines via ``logger.info(...)`` (lines 739?742 and 1230?1233) ? Python ``logging`` goes to **`**stderr**`** by default. But the native runner's live loop reads **`**only `proc.stdout`**`** (native_wasb_runner.py:385), capturing stderr only at the end via ``communicate()``. So the progress lines never reach the parser and ``_progress`` never fires. Let me confirm the exact emission format, the two code paths (CPU vs nvdec), and how logging is configured in ``wasb_infer.py``.

23. user (2026-07-09T06:26:22.810Z)

```

49     import argparse
50     import csv
51     import glob
52     import json
53     import logging
54     import os
55     import os.path as osp
56     from collections import defaultdict
57     from contextlib import contextmanager
58
59     logger = logging.getLogger(__name__)
60
61     CACHE_VERSION = 1
62     _DET_FILE = "det_raw.jsonl"
63     _MANIFEST_FILE = "manifest.json"
64
65     # Pixel-decode backends (#506). "cpu" = the historical cv2 paths (disk frames or
66     # --stream-video); "nvdec_resident" = GPU-resident NVDEC decode + affine grid_sample
67     # resize + normalize (torch 2.x + PyNvVideoCodec, CUDA-only).
68     DECODE_BACKENDS = ("cpu", "nvdec_resident")
69
70     # ImageNet statistics ? MUST match WASB's build_img_transforms Normalize (the GPU-
71     resident
72     # frontend normalizes on-device instead of through torchvision's transform).
73     IMAGENET_MEAN = (0.485, 0.456, 0.406)
74     IMAGENET_STD = (0.229, 0.224, 0.225)
75
76     def _build_cfg(weights: str, sport: str, device: str = "cuda"):
77         """Compose the WASB eval config via Hydra, overriding model/weights/device.
78

```

24. user (2026-07-09T06:26:23.663Z)

```

720         for ie in sorted(batch_results[ib].keys()):
721             p = img_paths[ie][ib]
722             fid = _frame_id(p)
723             plain = [_det_plain(d) for d in batch_results[ib][ie]]
724             frames_payload.append([fid, plain])
725             det_by_fid[fid].extend(plain)
726         pending.append({"w": win_idx, "f": frames_payload})
727         win_idx += 1

```

```

728         done += nb
729         if len(pending) >= flush_n or done >= windows_total:
730             _append_windows(det_f, pending)
731             pending = []
732             base_manifest["windows_done"] = done
733             write_manifest(cache_dir, base_manifest)
734         if (bi + 1) % log_every == 0 or done >= windows_total:
735             el = time.time() - t0
736             new_done = done - windows_done
737             rate = new_done / el if el > 0 else 0.0
738             eta = (windows_total - done) / rate if rate > 0 else 0.0
739             logger.info(
740                 f"{done}/{windows_total} "
741                 f"({100 * done // max(1, windows_total)}%) | {rate:.1f}
win/s | "
742                 f"elapsed {el:.0f}s | ETA {eta:.0f}s"
743             )
744         if pending:
745             _append_windows(det_f, pending)
746             base_manifest["windows_done"] = done
747             write_manifest(cache_dir, base_manifest)
748         finally:
749             det_f.close()
750             logger.info(f"detection done in {time.time() - t0:.0f}s; running tracker...")
751     else:
752         logger.info("detector cache already complete -> tracker-only re-run")
753
754     # ---- tracker re-run over the full ordered cache (cheap, deterministic) ---- #
755     return _tracker_pass_and_write(cfg, det_by_fid, out_csv)
756
757
758 def _tracker_pass_and_write(cfg, det_by_fid, out_csv: str) -> str:
759     """Tracker re-run over the full ordered detection cache ? trajectory CSV.
760
761     Shared verbatim by every decode backend: the tracker consumes the same per-frame
762     candidate lists regardless of how the pixels were decoded, so keeping this single
763     keeps the output contract single.
764
765     NB: this is WASB's OWN local `trackers` module (from its repo, on sys.path inside
766     the
767     `wasb` conda env), NOT the Roboflow `trackers` package. The `type: ignore`
768     suppresses
769     the mypy false positive where the main env resolves `trackers` to the Roboflow
770     package
771     (which has no `build_tracker`). There is NO runtime collision ? this runs in a
772     separate
773     subprocess/conda env. Do NOT "clean up" the ignore without removing the Roboflow
774     dep.
775
776     """
777     from trackers import build_tracker # type: ignore[attr-defined]
778
779     _ensure_np_inf() # OnlineTracker reads np.Inf (removed in NumPy 2.0)
780     tracker = build_tracker(cfg)

```

25. assistant (2026-07-09T06:26:33.415Z)

Confirmed: progress format `{done}/{windows\_total} ({pct}%) | ... win/s | elapsed ...` matches the regex ? the emission path is the problem, not the pattern. Let me confirm the `main()` logging config (does it route to stdout or stderr?) and that the nvdec branch emits identically.

26. user (2026-07-09T06:26:35.921Z)

```

1215         pending.append({"w": win_idx, "f": frames_payload})
1216         win_idx += 1
1217         done += nb

```

```

1218         n_batches += 1
1219         batch_fids, batch_imgs = [], []
1220         if len(pending) >= flush_n or done >= windows_total:
1221             _append_windows(det_f, pending)
1222             pending = []
1223             base_manifest["windows_done"] = done
1224             write_manifest(cache_dir, base_manifest)
1225         if n_batches % log_every == 0 or done >= windows_total:
1226             el = time.time() - t0
1227             new_done = done - windows_done
1228             rate = new_done / el if el > 0 else 0.0
1229             eta = (windows_total - done) / rate if rate > 0 else 0.0
1230             logger.info(
1231                 f"{done}/{windows_total} "
1232                 f"({100 * done // max(1, windows_total)}%) | {rate:.1f}
win/s | "
1233                 f"elapsed {el:.0f}s | ETA {eta:.0f}s"
1234             )
1235         if pending: # safety net; the done>=windows_total branch already
flushed
1236             _append_windows(det_f, pending)
1237             base_manifest["windows_done"] = done
1238             write_manifest(cache_dir, base_manifest)
1239         finally:
1240             det_f.close()
1241             logger.info(f"detection done in {time.time() - t0:.0f}s; running
tracker...")
1242         else:
1243             logger.info("detector cache already complete -> tracker-only re-run")
1244         finally:

```

27. user (2026-07-09T06:26:36.566Z)

```

1250 def main():
1251     ap = argparse.ArgumentParser()
1252     ap.add_argument("--frames_dir", help="folder of pre-extracted frames")
1253     ap.add_argument(
1254         "--video", help="video file; frames are extracted internally (in WSL)"
1255     )
1256     ap.add_argument(
1257         "--frames_out_dir", help="where to extract frames when --video is used"
1258     )
1259     ap.add_argument(
1260         "--stream-video",
1261         action="store_true",
1262         help="FAST PATH: decode --video on the fly and feed frames straight to "
1263         "the detector (no PNG extraction). Output-equivalent to the disk path.",
1264     )
1265     ap.add_argument(
1266         "--decode-backend",
1267         default="cpu",
1268         choices=list(DECODE_BACKENDS),
1269         help="cpu (default) = the historical cv2 decode paths (disk frames or "
1270         "--stream-video). nvdec_resident = GPU-RESIDENT decode+resize+normalize "
1271         "(#506): PyNvVideoCodec NVDEC + cv2-faithful affine grid_sample, CUDA-only, "
1272         "~1.7x on an L4. Requires --video; detections cache separately from cpu "
1273         "(trajectory-level, not byte-level, equivalence).",
1274     )
1275     ap.add_argument("--weights", required=True)
1276     ap.add_argument("--out", required=True)
1277     ap.add_argument("--sport", default="badminton")
1278     ap.add_argument(
1279         "--device",
1280         default="cuda",
1281         choices=["cuda", "mps", "cpu"],

```

```

1282         help="torch device for the detector (default cuda = the WSL parity path)",
1283     )
1284     ap.add_argument(
1285         "--limit", type=int, default=0, help="cap #frames (for quick tests)"
1286     )
1287     ap.add_argument("--batch-size", type=int, default=8)
1288     ap.add_argument("--num-workers", type=int, default=4)
1289     ap.add_argument(
1290         "--prefetch-batches",
1291         type=int,
1292         default=0,
1293         help="streaming mode: overlap CPU batch assembly with GPU inference via a "
1294         "bounded producer thread of this queue depth (0 = off, the parity default)",
1295     )
1296     ap.add_argument(
1297         "--log-every-batches", type=int, default=50, help="progress log cadence"
1298     )
1299     ap.add_argument(
1300         "--cache-dir",
1301         default=None,
1302         help="resumable detector cache dir (default: <out>_wasbcache next to --out)",
1303     )
1304     ap.add_argument(
1305         "--flush-every",
1306         type=int,
1307         default=1000,
1308         help="fsync the detector cache every N windows (bounds reboot loss)",
1309     )
1310     ap.add_argument(
1311         "--fresh", action="store_true", help="ignore any existing cache and start over"
1312     )
1313     args = ap.parse_args()
1314
1315     if args.decode_backend == "nvdec_resident":
1316         # GPU-resident branch FIRST: it decodes the video itself (NVDEC), so neither the
1317         # --stream-video cv2 path nor the extract-frames fallback below may run.
1318         if not args.video:
1319             raise SystemExit("--decode-backend nvdec_resident requires --video")
1320         if args.device != "cuda":
1321             raise SystemExit(
1322                 "--decode-backend nvdec_resident requires --device cuda "
1323                 "(NVDEC decode runs on the GPU)"
1324             )
1325         run_video_nvdec_resident(
1326             args.video,
1327             args.weights,
1328             args.out,
1329             args.sport,
1330             args.limit,
1331             batch_size=args.batch_size,
1332             log_every_batches=args.log_every_batches,
1333             cache_dir=args.cache_dir,
1334             flush_every=args.flush_every,
1335             fresh=args.fresh,
1336         )
1337         return
1338
1339     if args.stream_video:
1340         if not args.video:
1341             raise SystemExit("--stream-video requires --video")
1342         run_video_streaming(
1343             args.video,
1344             args.weights,
1345             args.out,
1346             args.sport,

```



```

1347         args.limit,
1348         batch_size=args.batch_size,
1349         log_every_batches=args.log_every_batches,
1350         cache_dir=args.cache_dir,
1351         flush_every=args.flush_every,
1352         fresh=args.fresh,
1353         device=args.device,
1354         prefetch_batches=args.prefetch_batches,
1355     )
1356     return
1357
1358     frames_dir = args.frames_dir
1359     if args.video:
1360         frames_dir = extract_frames(
1361             args.video, args.frames_out_dir or (args.video + "_frames")
1362         )
1363     if not frames_dir:
1364         raise SystemExit("provide --frames_dir or --video")
1365     run(
1366         frames_dir,
1367         args.weights,
1368         args.out,
1369         args.sport,
1370         args.limit,
1371         batch_size=args.batch_size,
1372         num_workers=args.num_workers,
1373         log_every_batches=args.log_every_batches,
1374         cache_dir=args.cache_dir,
1375         flush_every=args.flush_every,
1376         fresh=args.fresh,
1377         device=args.device,
1378     )
1379
1380
1381     if __name__ == "__main__":
1382         main()
1383

```

28. assistant (2026-07-09T06:27:34.619Z)

Now let me check a decisive detail: does `wasb\_infer.py` configure logging at all? If it doesn't, `logger.info(...)` sits below the root logger's default WARNING level and is ****suppressed entirely**** (not even reaching stderr) ? which changes the fix. Let me verify.

29. user (2026-07-09T06:27:37.227Z)

No matches found

30. user (2026-07-09T06:27:38.032Z)

```

1040         f"expected 3-channel RGB (is output_color_type=RGB supported here?)"
1041     )
1042     return t
1043
1044     def read_chw(self, index: int):
1045         """Decoded frame ``index`` ? the normalized ``[3, inp_h, inp_w]`` float32 CUDA
1046         tensor the detector expects (warpAffine-faithful resize + ImageNet
1047         normalize)."""
1048         import torch.nn.functional as F
1049
1050         hwc = self._decode_rgb_hwc(index)
1051         chw = hwc.permute(2, 0, 1).unsqueeze(0).float() / 255.0
1052         # padding_mode="zeros" = cv2.warpAffine's BORDER_CONSTANT(0) black fill ? see
1053         # build_affine_grid's docstring (only non-16:9 sources ever sample the border).

```

```

1053         warped = F.grid_sample(
1054             chw, self._grid, mode="bilinear", align_corners=False, padding_mode="zeros"
1055         ).squeeze(0)
1056         return (warped - self._mean) / self._std
1057
1058     def release(self):
1059         self._dec = None
1060
1061
1062     def run_video_nvdec_resident(
1063         video_path: str,
1064         weights: str,
1065         out_csv: str,
1066         sport: str = "badminton",
1067         limit: int = 0,
1068         batch_size: int = 8,
1069         log_every_batches: int = 50,

```

31. assistant (2026-07-09T06:29:51.949Z)

Plumbing check + score source. The grep already confirmed ``process_video(progress=...)`` forwards to ``run_predict(progress=progress)`` (trajectory_hybrid.py:300), so the callback IS wired ? GAP 3 is purely emission/capture. Two compounding bugs there: (1) ``wasb_infer.py`` **never** configures logging, so ``logger.info`` progress sits below root WARNING and is largely suppressed; (2) even when emitted it goes to **stderr**, but the runner reads **stdout only** (and draining stdout while stderr's 64 KB pipe fills risks a deadlock). Now let me confirm the progress plumbing has no early-out, and find where validators put a numeric "score" in ``details``.

32. user (2026-07-09T06:29:53.958Z)

```

170         Persisted candidates are NOT affected ? they remain the full superset for
offline
171         re-scoring. """
172         return list(range(len(windows)))
173
174     def process_video( # type: ignore[override]
175         self,
176         video_id: str,
177         video_path: str,
178         player_pool: Optional[List[str]] = None,
179         max_frames: Optional[int] = None,
180         progress: Optional[Callable[[str], None]] = None,
181     ) -> "Union[List[Dict[str, Any]], Failure]":
182         # Result contract (the ABC declares Success|Failure). Anything that means the
run could
183         # NOT actually complete ? an unrunnable config (unknown sport/provider, missing
API key),
184         # an unreadable video, an unavailable detector (native-not-supported), the env
healthcheck
185         # not ready, or run_predict timing out / aborting with no trajectory ? returns a
`Failure`
186         # so it is NOT confused with a genuine 0-rally video. A LEGITIMATE empty match
(a trajectory
187         # was produced but no candidate/AI-confirmed rallies) still returns a bare
`[]`. Every
188         # caller (worker.cmd_infer, main.py, orchestration) branches on
isinstance(result, Failure).
189         label = self.display_label
190         # --- Sport profile + AI provider wiring (identical across segmenters) ---
191         sport_name = self.config.sport
192         try:
193             sport_profile = sport_registry.get(sport_name)
194         except KeyError as e:
195             logger.error(str(e))
196             # Unrunnable config, not a 0-rally match (mirrors the gemini segmenter).

```

```
197         return Failure(str(e))
198
199         idx_cfg = self.config.indexing
```

33. user (2026-07-09T06:29:55.393Z)

<persisted-output>

Output too large (29.2KB). Full output saved to: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\663947f-2127-4383-9905-650340ff9863\tool-results\toolu_012XtzFntSb45bw2HyhK5Chq.txt

Preview (first 2KB):

```
backend\pipeline\validators\base.py:26:class ValidationResult(BaseModel):
backend\pipeline\validators\base.py-27-     validator_name: str
backend\pipeline\validators\base.py-28-     passed: bool
backend\pipeline\validators\base.py-29-     message: str
backend\pipeline\validators\base.py-30-     details: dict[str, Any] = {}
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-31-
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-32-
--
backend\pipeline\validators\base.py:148:         ValidationResult(
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-149-                                     validator_name=val.name,
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-150-                                     passed=False,
backend\pipeline\validators\base.py-151-                                     message=f"Validator threw an
unexpected error: {str(e)}",
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-152-                                     )
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-153-                                     )
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-154-         return results, context
--
backend\pipeline\validators\blur.py:39:         return ValidationResult(
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-40-                                     validator_name=self.name,
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-41-                                     passed=False,
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-42-                                     message="Failed to open video
file for blur analysis.",
...
</persisted-output>
```

34. assistant (2026-07-09T06:30:02.089Z)

Let me read the full validator-details survey to understand the `details` structure and find the score keys.

35. user (2026-07-09T06:30:03.904Z)

```
1     backend\pipeline\validators\base.py:26:class ValidationResult(BaseModel):
2     backend\pipeline\validators\base.py-27-         validator_name: str
3     backend\pipeline\validators\base.py-28-         passed: bool
4     backend\pipeline\validators\base.py-29-         message: str
5     backend\pipeline\validators\base.py-30-         details: dict[str, Any] = {}
6     C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-31-
7     C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-32-
8     --
```

```

9         backend\pipeline\validators\base.py:148:                               ValidationResult(
10         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-149-                               validator_name=val.name,
11         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-150-                               passed=False,
12         backend\pipeline\validators\base.py-151-                               message=f"Validator
threw an unexpected error: {str(e)}",
13         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-152-                               )
14         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-153-                               )
15         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-154-                               return results, context
16         --
17         backend\pipeline\validators\blur.py:39:                               return ValidationResult(
18         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-40-                               validator_name=self.name,
19         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-41-                               passed=False,
20         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-42-                               message="Failed to open video
file for blur analysis.",
21         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-43-                               )
22         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-44-
23         backend\pipeline\validators\blur.py-45-                               if meta.total_frames <= 0:
24         backend\pipeline\validators\blur.py-46:                               return ValidationResult(
25         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-47-                               validator_name=self.name,
passed=False, message="Invalid frame count."
26         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-48-                               )
27         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-49-
28         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-50-                               variances = state["variances"]
29         backend\pipeline\validators\blur.py-51-                               if not variances:
30         backend\pipeline\validators\blur.py-52:                               return ValidationResult(
31         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-53-                               validator_name=self.name,
32         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-54-                               passed=False,
33         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-55-                               message="Could not read frames to
perform blur verification.",
34         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-56-                               )
35         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-57-
36         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-58-                               avg_variance = sum(variances) /
len(variances)
37         --
38         backend\pipeline\validators\blur.py:71:                               return ValidationResult(
39         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-72-                               validator_name=self.name,
40         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-73-                               passed=False,
41         backend\pipeline\validators\blur.py-74-                               message=f"Video quality rejected:
Focus check failed. Average sharpness score of {round(avg_variance, 1)} is below the minimum
required threshold of {min_blur_threshold}.",
42         backend\pipeline\validators\blur.py:75:                               details=details,
43         C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-76-                               )

```

```

44     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-77-
45     backend\pipeline\validators\blur.py:78:         return ValidationResult(
46     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-79-         validator_name=self.name,
47     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-80-         passed=True,
48     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-81-         message="Video sharpness meets focus
requirements.",
49     backend\pipeline\validators\blur.py:82:         details=details,
50     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py-83-     )
51     --
52     backend\pipeline\validators\format.py:27:         return ValidationResult(
53     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-28-         validator_name=self.name,
54     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-29-         passed=False,
55     backend\pipeline\validators\format.py-30-         message=f"File not found:
{video_path}",
56     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-31-     )
57     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-32-
58     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-33-         cmd = [
59     --
60     backend\pipeline\validators\format.py:54:         return ValidationResult(
61     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-55-         validator_name=self.name,
62     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-56-         passed=False,
63     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-57-         message="Failed to retrieve
valid video streams from the file.",
64     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-58-     )
65     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-59-
66     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-60-         # Store the metadata in context for
subsequent validators to utilize
67     --
68     backend\pipeline\validators\format.py:75:         return ValidationResult(
69     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-76-         validator_name=self.name,
70     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-77-         passed=False,
71     backend\pipeline\validators\format.py-78-         message=f"Invalid format:
Found '{fmt.get('format_name')}', but MP4 is required.",
72     backend\pipeline\validators\format.py:79:         details=details,
73     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-80-     )
74     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-81-
75     backend\pipeline\validators\format.py:82:         return ValidationResult(
76     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-83-         validator_name=self.name,
77     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-84-         passed=True,
78     C:\Users\avidu\Projects\kheelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-85-         message="Video format is valid
MP4 container.",
79     backend\pipeline\validators\format.py:86:         details=details,

```

```

80     C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-87-        )
81     C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-88-
82     backend\pipeline\validators\format.py-89-        except subprocess.CalledProcessError as
e:
83         backend\pipeline\validators\format.py:90:            return ValidationResult(
84         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-91-            validator_name=self.name,
85         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-92-            passed=False,
86         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-93-            message="Video stream is
corrupted or unreadable by ffprobe.",
87         backend\pipeline\validators\format.py:94:            details={"error": e.stderr},
88         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-95-        )
89         backend\pipeline\validators\format.py-96-        except subprocess.TimeoutExpired:
90         backend\pipeline\validators\format.py:97:            return ValidationResult(
91         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-98-            validator_name=self.name,
92         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-99-            passed=False,
93         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-100-            message=f"ffprobe timed out
after {timeout_label(ffprobe_timeout_sec())}.",
94         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-101-        )
95         backend\pipeline\validators\format.py-102-        except Exception as e:
96         backend\pipeline\validators\format.py:103:            return ValidationResult(
97         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-104-            validator_name=self.name,
98         C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-105-            passed=False,
99         backend\pipeline\validators\format.py-106-            message=f"Unexpected error
running ffprobe: {str(e)}",
100        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-107-        )
101        --
102        backend\pipeline\validators\pts_continuity.py:46:            return
ValidationResult(
103        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-47-            validator_name=self.name,
104        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-48-            passed=False,
105        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-49-            message="No packet
timestamps could be extracted from the video stream.",
106        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-50-        )
107        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-51-
108        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-52-            timestamps = []
109        --
110        backend\pipeline\validators\pts_continuity.py:66:            return
ValidationResult(
111        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-67-            validator_name=self.name,
112        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-68-            passed=True,
113        C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-69-            message="Video is

```

```

too short to perform a PTS continuity check.",
114     backend\pipeline\validators\pts_continuity.py:70:
details={"keyframes_count": len(timestamps)},
115     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-71-
116     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-72-
117     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-73-
118     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-74-
119     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-75-
120     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-76-
121     --
122     backend\pipeline\validators\pts_continuity.py:98:
ValidationResult(
123     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-99-
validator_name=self.name,
124     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-100-
125     backend\pipeline\validators\pts_continuity.py-101-
integrity compromise: Backward timeline jumps (splicing) detected.",
126     backend\pipeline\validators\pts_continuity.py:102:
127     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-103-
128     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-104-
129     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-105-
config.validation.pts_max_gap_seconds
130     backend\pipeline\validators\pts_continuity.py-106-
max_gap_threshold:
131     backend\pipeline\validators\pts_continuity.py:107:
ValidationResult(
132     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-108-
validator_name=self.name,
133     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-109-
134     backend\pipeline\validators\pts_continuity.py-110-
integrity compromise: Detected a timestamp gap of {round(max_gap, 2)}s, indicating edited or
missing content.",
135     backend\pipeline\validators\pts_continuity.py:111:
136     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-112-
137     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-113-
138     backend\pipeline\validators\pts_continuity.py:114:
139     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-115-
validator_name=self.name,
140     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-116-
141     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-117-
is continuous with no abnormal gaps or backward jumps.",
142     backend\pipeline\validators\pts_continuity.py:118:
143     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-119-
144     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-120-
145     backend\pipeline\validators\pts_continuity.py-121-
except

)

# Analyze gaps
max_gap = 0.0
backward_jumps = 0
gaps = []

return

passed=False,
message="Video

details=details,

)

max_gap_threshold =

if max_gap >

return

passed=False,
message=f"Video

details=details,

)

return ValidationResult(

passed=True,
message="PTS timeline

details=details,

)

except

```

```

subprocess.CalledProcessError as e:
146     backend\pipeline\validators\pts_continuity.py:122:         return ValidationResult(
147     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-123-
validator_name=self.name,
148     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-124-         passed=False,
149     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-125-         message="Failed to
execute ffprobe PTS analysis.",
150     backend\pipeline\validators\pts_continuity.py:126:         details={"error":
e.stderr},
151     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-127-         )
152     backend\pipeline\validators\pts_continuity.py-128-         except
subprocess.TimeoutExpired:
153     backend\pipeline\validators\pts_continuity.py:129:         return ValidationResult(
154     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-130-
validator_name=self.name,
155     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-131-         passed=False,
156     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-132-         message=f"ffprobe PTS
analysis timed out after {timeout_label(ffprobe_timeout_sec())}.",
157     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-133-         )
158     backend\pipeline\validators\pts_continuity.py-134-         except Exception as e:
159     backend\pipeline\validators\pts_continuity.py:135:         return ValidationResult(
160     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-136-
validator_name=self.name,
161     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-137-         passed=False,
162     backend\pipeline\validators\pts_continuity.py-138-         message=f"Unexpected
error in PTS validation: {str(e)}",
163     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-139-         )
164     --
165     backend\pipeline\validators\relevance.py:80:         return ValidationResult(
166     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-81-         validator_name=self.name,
167     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-82-         passed=False,
168     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-83-         message="Failed to open
video file for relevance validation.",
169     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-84-         )
170     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-85-
171     backend\pipeline\validators\relevance.py-86-         if meta.total_frames <= 0:
172     backend\pipeline\validators\relevance.py:87:         return ValidationResult(
173     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-88-         validator_name=self.name,
passed=False, message="Invalid frame count."
174     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-89-         )
175     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-90-
176     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-91-         green_ratios = state["green_ratios"]
177     backend\pipeline\validators\relevance.py-92-         if not green_ratios:
178     backend\pipeline\validators\relevance.py:93:         return ValidationResult(
179     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-

```



```

indexer\backend\pipeline\validators\relevance.py-94-           validator_name=self.name,
180     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-95-           passed=False,
181     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-96-           message="Failed to extract
frames for relevance checking.",
182     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-97-           )
183     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-98-
184     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-99-           sport_name = config.sport
185     --
186     backend\pipeline\validators\relevance.py:112:           return ValidationResult(
187     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-113-           validator_name=self.name,
188     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-114-           passed=False,
189     backend\pipeline\validators\relevance.py-115-           message=f"Video rejected:
Relevance check failed. The scene does not contain a recognizable {sport_name} court (court area
found: {round(avg_green_ratio * 100, 2)}% vs min required {round(min_court_ratio * 100, 2)}%).",
190     backend\pipeline\validators\relevance.py:116:           details=details,
191     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-117-           )
192     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-118-
193     backend\pipeline\validators\relevance.py:119:           return ValidationResult(
194     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-120-           validator_name=self.name,
195     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-121-           passed=True,
196     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-122-           message=f"Video identified as
relevant to {sport_name} ({sport_name} court structures found).",
197     backend\pipeline\validators\relevance.py:123:           details=details,
198     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-124-           )
199     --
200     backend\pipeline\validators\resolution_fps.py:68:           return ValidationResult(
201     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-69-           validator_name=self.name,
202     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-70-           passed=False,
203     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-71-           message="Video metadata
is missing. Run format_check validator first.",
204     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-72-           )
205     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-73-
206     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-74-           stream = meta["streams"][0]
207     --
208     backend\pipeline\validators\resolution_fps.py:95:           return ValidationResult(
209     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-96-           validator_name=self.name,
210     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-97-           passed=False,
211     backend\pipeline\validators\resolution_fps.py-98-           message=f"Low
resolution: Found {width}x{height}, but minimum is {min_w}x{min_h}.",
212     backend\pipeline\validators\resolution_fps.py:99:           details=details,
213     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-100-           )

```

```

214     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-101-
215     backend\pipeline\validators\resolution_fps.py-102-         if fps < min_fps:
216     backend\pipeline\validators\resolution_fps.py-103:         return ValidationResult(
217     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-104-
validator_name=self.name,
218     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-105-         passed=False,
219     backend\pipeline\validators\resolution_fps.py-106-         message=f"Low frame
rate: Found {round(fps, 2)} FPS, but minimum is {min_fps} FPS.",
220     backend\pipeline\validators\resolution_fps.py-107:         details=details,
221     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-108-         )
222     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-109-
223     backend\pipeline\validators\resolution_fps.py-110:         return ValidationResult(
224     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-111-         validator_name=self.name,
225     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-112-         passed=True,
226     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-113-         message="Video resolution
and frame rate meet minimum standards.",
227     backend\pipeline\validators\resolution_fps.py-114:         details=details,
228     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-115-         )
229     --
230     backend\pipeline\validators\scene_cut.py-62:         return ValidationResult(
231     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-63-         validator_name=self.name,
232     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-64-         passed=False,
233     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-65-         message="Failed to open
video file using OpenCV.",
234     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-66-         )
235     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-67-
236     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-68-         fps = meta.fps
237     --
238     backend\pipeline\validators\scene_cut.py-73:         return ValidationResult(
239     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-74-         validator_name=self.name,
240     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-75-         passed=False,
241     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-76-         message="Invalid video
duration.",
242     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-77-         )
243     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-78-
244     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-79-         sample_interval = max(1, int(fps /
2))
245     --
246     backend\pipeline\validators\scene_cut.py-99:         return ValidationResult(
247     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-100-         validator_name=self.name,
248     C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-101-         passed=False,
249     backend\pipeline\validators\scene_cut.py-102-         message=f"Video splicing

```

```
compromise: Detected {cuts_detected} cuts ({round(cuts_per_minute, 2)} cuts/min), which exceeds
the continuous raw file limit of {max_cuts_rate} cuts/min.",
250     backend\pipeline\validators\scene_cut.py:103:                 details=details,
251
252     [Showing results with pagination = limit: 250]
```

36. assistant (2026-07-09T06:30:31.948Z)

The validators each put their own numeric scores in `details` (blur?sharpness, relevance?court ratio, scene_cut?cuts/min, etc.) ? there's no uniform `score` field. Note `[r.model\_dump() for r in val\_results]` (already stored to the ephemeral DB at line 306) yields exactly `{validator\_name, passed, message, details}` ? the clean thing to also put in report.json. Let me read the `details` dicts to design a robust compact score hint for the log line.

37. user (2026-07-09T06:30:32.958Z)

```
1     from typing import Any, List
2
3     import cv2
4
5     from backend.pipeline.validators.base import FrameSamplingValidator, ValidationResult
6     from backend.pipeline.validators.frame_sampler import SampleMeta
7
8     # Number of frames sampled evenly across the whole video.
9     _NUM_SAMPLES = 15
10
11
12     class BlurValidator(FrameSamplingValidator):
13         @property
14         def name(self) -> str:
15             return "blur_check"
16
17         @property
18         def description(self) -> str:
19             return "Samples frames across the video and checks if the footage is excessively
blurry or out of focus."
20
21         def sample_indices(self, total_frames: int, fps: float) -> List[int]:
22             # Sample _NUM_SAMPLES frames evenly distributed across the video.
23             step = max(1, total_frames // _NUM_SAMPLES)
24             return [i * step for i in range(_NUM_SAMPLES)]
25
26         def _init_state(self, config: Any) -> dict[str, Any]:
27             return {"variances": []}
28
29         def _consume(self, state: dict[str, Any], frame_idx: int, frame: Any) -> None:
30             gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
31             # Calculate the Laplacian variance (standard focus metric)
32             laplacian_var = cv2.Laplacian(gray, cv2.CV_64F).var()
33             state["variances"].append(laplacian_var)
34
35         def _finalize(
36             self, state: dict[str, Any], meta: SampleMeta, config: Any, context: dict
37         ) -> ValidationResult:
38             if not meta.opened:
39                 return ValidationResult(
40                     validator_name=self.name,
41                     passed=False,
42                     message="Failed to open video file for blur analysis.",
43                 )
44
45             if meta.total_frames <= 0:
46                 return ValidationResult(
47                     validator_name=self.name, passed=False, message="Invalid frame count."
48                 )
```

```

49
50     variances = state["variances"]
51     if not variances:
52         return ValidationResult(
53             validator_name=self.name,
54             passed=False,
55             message="Could not read frames to perform blur verification.",
56         )
57
58     avg_variance = sum(variances) / len(variances)
59
60     details = {
61         "avg_laplacian_variance": round(avg_variance, 2),
62         "min_laplacian_variance": round(min(variances), 2),
63         "max_laplacian_variance": round(max(variances), 2),
64         "sampled_frames_count": len(variances),
65     }
66
67     val_cfg = config.validation
68     min_blur_threshold = val_cfg.min_blur_threshold
69
70     if avg_variance < min_blur_threshold:
71         return ValidationResult(
72             validator_name=self.name,
73             passed=False,
74             message=f"Video quality rejected: Focus check failed. Average sharpness
score of {round(avg_variance, 1)} is below the minimum required threshold of
{min_blur_threshold}.",
75             details=details,
76         )
77
78     return ValidationResult(
79         validator_name=self.name,
80         passed=True,
81         message="Video sharpness meets focus requirements.",
82         details=details,
83     )
84

```

38. user (2026-07-09T06:30:34.784Z)

```

backend\pipeline\validators\blur.py:60:         details = {
backend\pipeline\validators\blur.py:61-         "avg_laplacian_variance": round(avg_variance,
2),
backend\pipeline\validators\blur.py:62-         "min_laplacian_variance":
round(min(variances), 2),
backend\pipeline\validators\blur.py:63-         "max_laplacian_variance":
round(max(variances), 2),
backend\pipeline\validators\blur.py:64-         "sampled_frames_count": len(variances),
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py:65-     }
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py:66-
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\blur.py:67-     val_cfg = config.validation
--
backend\pipeline\validators\pts_continuity.py:89:         details = {
backend\pipeline\validators\pts_continuity.py:90-         "keyframes_count":
len(timestamps),
backend\pipeline\validators\pts_continuity.py:91-         "max_keyframe_gap_seconds":
round(max_gap, 2),
backend\pipeline\validators\pts_continuity.py:92-         "backward_jumps_detected":
backward_jumps,
backend\pipeline\validators\pts_continuity.py:93-         "first_keyframe":
timestamps[0],

```

```

backend\pipeline\validators\pts_continuity.py-94-         "last_keyframe":
timestamps[-1],
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-95-         }
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\pts_continuity.py-96-
--
backend\pipeline\validators\relevance.py:103:         details = {
backend\pipeline\validators\relevance.py-104-             "avg_green_court_ratio":
round(avg_green_ratio, 4),
backend\pipeline\validators\relevance.py-105-             "max_green_court_ratio":
round(max(green_ratios), 4),
backend\pipeline\validators\relevance.py-106-             "min_green_court_ratio":
round(min(green_ratios), 4),
backend\pipeline\validators\relevance.py-107-             "green_threshold": min_court_ratio,
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-108-         }
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-109-
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\relevance.py-110-         # If a certain amount of court
color is visible, consider it relevant
--
backend\pipeline\validators\format.py:66:         details = {
backend\pipeline\validators\format.py-67-             "format_name": fmt.get("format_name"),
backend\pipeline\validators\format.py-68-             "duration": float(fmt.get("duration",
0)),
backend\pipeline\validators\format.py-69-             "size_bytes": int(fmt.get("size", 0)),
backend\pipeline\validators\format.py-70-             "codec_name": stream.get("codec_name"),
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-71-         }
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-72-
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\format.py-73-         # Check if container name includes
mp4
--
backend\pipeline\validators\base.py:30:     details: dict[str, Any] = {}
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-31-
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-32-
backend\pipeline\validators\base.py-33-class VideoValidator(ABC):
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-34-     @property
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-35-     @abstractmethod
backend\pipeline\validators\base.py-36-     def name(self) -> str:
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\base.py-37-         """Unique identifier for this
validator."""
--
backend\pipeline\validators\resolution_fps.py:86:         details = {"width": width, "height":
height, "fps": round(fps, 2)}
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-87-
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-88-         # Check bounds from typed
config (AppConfig.validation)
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-89-         val_cfg = config.validation
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-90-         min_w =
val_cfg.min_resolution_width
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-

```

```

indexer\backend\pipeline\validators\resolution_fps.py-91-         min_h =
val_cfg.min_resolution_height
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-92-         min_fps = val_cfg.min_fps
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\resolution_fps.py-93-
--
backend\pipeline\validators\scene_cut.py:89:         details = {
backend\pipeline\validators\scene_cut.py-90-             "cuts_detected": cuts_detected,
backend\pipeline\validators\scene_cut.py-91-             "sampled_frames": sampled_count,
backend\pipeline\validators\scene_cut.py-92-             "cuts_per_minute":
round(cuts_per_minute, 3),
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-93-         }
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-94-
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-95-         val_cfg = config.validation
C:\Users\avidu\Projects\khehsutra-guru\badminton-highlight-
indexer\backend\pipeline\validators\scene_cut.py-96-         max_cuts_rate =
val_cfg.max_scene_cuts_per_minute

```

39. assistant (2026-07-09T06:32:08.697Z)

Now I have the full `details` picture ? each validator has its own metric keys (blur?`avg\_laplacian\_variance`, relevance?`avg\_green\_court\_ratio`, scene_cut?`cuts\_per\_minute`, etc.), no uniform `score`. For report.json I'll persist `[r.model\_dump() ?]` (same list already going to the DB ? carries `details` with all scores); for the log line I'll derive a compact score hint from known keys.

Let me look at the test infrastructure: the module-load-by-path pattern to mirror, and any existing tests for the worker infer path and the native runner's progress streaming (added by PR #546).

40. user (2026-07-09T06:32:10.700Z)

```

1      """Tests for scripts/gates/quality/nightly_regression.py ? the pure
summary/comparison/report core.
2
3      Fully offline: no real golden corpus, no trajectory I/O, no GPU/numpy. The comparison
core
4      operates on plain summary dicts, so a regression / over-seg-guard / corpus-change can be
5      exercised deterministically. The real-corpus scoring path (`run_corpus`) is a thin
wrapper
6      over ``rally_seg_eval.evaluate`` (tested in test_rally_seg_eval.py) and is exercised e2e
on the
7      owner's machine, not here.
8      """
9
10     import importlib.util
11     import os
12     import sys
13
14     import pytest
15
16     # scripts/ isn't a package ? load the module by path. Register in sys.modules BEFORE
exec so
17     # dataclasses can resolve the `from __future__ import annotations` string annotations.
18     _SCRIPT = os.path.join(
19         os.path.dirname(os.path.dirname(os.path.abspath(__file__))),
20         "scripts",
21         "gates",
22         "quality",
23         "nightly_regression.py",
24     )

```

```

25     _spec = importlib.util.spec_from_file_location("nightly_regression", _SCRIPT)
26     nr = importlib.util.module_from_spec(_spec)
27     sys.modules["nightly_regression"] = nr
28     _spec.loader.exec_module(nr)
29
30
31     # ----- #
32     # Builders
33     # ----- #
34     _PRESET = {
35         "name": "x",
36         "velocity_thresh": 0.02,
37         "merge_gap": 2.0,
38         "min_window_duration": 2.0,
39         "max_window_duration": 0.0,
40         "hard_cap": False,
41         "inactivity_end": 0.0,
42         "inactivity_rally_thresh": 0.08,
43         "inactivity_min_density": 0.34,
44     }
45
46
47     def _summary(aggregate, per_video, preset=None, min_crossings=0, n=None):
48         return {
49             "n": n if n is not None else len(per_video),
50             "preset": preset or dict(_PRESET),
51             "min_crossings": min_crossings,
52             "iou": 0.5,
53             "aggregate": dict(aggregate),
54             "per_video": {k: dict(v) for k, v in per_video.items()},
55         }
56
57
58     def _pv(tol_f1=0.3, split_count=2, merge_count=0, **extra):
59         d = {
60             "tol_f1": tol_f1,
61             "f1@0.5": 0.4,
62             "f1@0.25": 0.5,
63             "merge_rate": 0.1,
64             "split_count": split_count,
65             "merge_count": merge_count,
66             "num_pred": 10,
67             "num_gt": 10,
68         }
69         d.update(extra)
70         return d

```

41. user (2026-07-09T06:32:10.891Z)

```

tests\__init__.py
tests\test_hashing.py
tests\test_import_smoke.py
tests\test_reliability_hardening.py
tests\test_annotation_providers.py
tests\test_app_home.py
tests\test_automated_provider.py
tests\test_base.py
tests\test_bind_and_host.py
tests\test_calibrate_wasb.py
tests\test_calibration.py
tests\test_cleanup_caches.py
tests\test_court_gate.py
tests\test_db_upgrade_safety.py
tests\test_detector_keying.py
tests\test_device_context.py

```

tests\test_distill_local.py
tests\test_edge_auth.py
tests\test_eval_annotations.py
tests\test_eval_batch.py
tests\test_eval_classifier.py
tests\test_eval_harness.py
tests\test_eval_metrics.py
tests\test_eval_partial_labels.py
tests\test_eval_regression.py
tests\test_eval_stats.py
tests\test_eval_training.py
tests\test_experiment.py
tests\test_experiment_honest.py
tests\test_export_wasb_segments.py
tests\test_flywheel.py
tests\test_fusion_features.py
tests\test_fusion_hybrid.py
tests\test_gemini.py
tests\test_gemini_refine.py
tests\test_generate_highlights.py
tests\test_geometry.py
tests\test_golden_fixtures.py
tests\test_golden_fixtures_split.py
tests\test_golden_real_fixtures.py
tests\test_gt_loader.py
tests\test_hit_detector.py
tests\test_job_self_scheduling.py
tests\test_job_worker_module_seam.py
tests\test_leak_scan.py
tests\test_light_tier_no_torch.py
tests\test_logging_setup.py
tests\test_motion_segmenter.py
tests\test_output_base.py
tests\test_p2b_install.py
tests\test_per_user_eval.py
tests\test_per_user_gate.py
tests\test_per_user_regression.py
tests\test_person_detector.py
tests\test_personalization_serving.py
tests\test_profile_parity.py
tests\test_promotion.py
tests\test_providers.py
tests\test_provider_gemini.py
tests\test_rally_seg_eval_report.py
tests\test_rally_seq_proto.py
tests\test_rally_slicer.py
tests\test_rally_summary.py
tests\test_redact.py
tests\test_reingest_safety.py
tests\test_reporting.py
tests\test_reporting_harvest.py
tests\test_reporting_spec_faq.py
tests\test_results.py
tests\test_run_signals.py
tests\test_run_telemetry.py
tests\test_score_calibration.py
tests\test_segmentation_metrics.py
tests\test_server_entrypoint.py
tests\test_serve_contrast.py
tests\test_single_instance.py
tests\test_splits.py
tests\test_sport_profiles.py
tests\test_stub_runner.py
tests\test_tolerance_metrics.py
tests\test_training_gen0.py


```
tests\test_validators_base.py
tests\test_variant_tuning.py
tests\test_watermark_i18n.py
tests\test_workspace.py
tests\test_wsl_tilde_quoting.py
tests\test_rally_gate.py
tests\test_rally_seg_eval.py
tests\test_windowing.py
tests\test_check_md_links.py
tests\test_freespace.py
tests\test_storage_capabilities.py
tests\test_tracknet_runner.py
tests\test_wasb_runner.py
tests\test_api_annotations.py
tests\test_annotation_proxy.py
tests\test_ffmpeg_resolver.py
tests\test_ffmpeg_timeouts.py
tests\test_video.py
tests\test_api_assets.py
```

(Showing 100 of 153 matching files; 53 more are not listed. Narrow the pattern or path to see the rest.)

42. assistant (2026-07-09T06:32:21.143Z)

Let me find the existing tests that touch the worker infer path, the native runner, and the #534 progress/telemetry work, so I extend rather than duplicate.

43. user (2026-07-09T06:32:24.046Z)

```
tests\test_native_wasb_runner.py
```

44. user (2026-07-09T06:32:25.294Z)

```
Found 17 files
tests\test_worker.py
tests\test_native_wasb_runner.py
tests\test_config.py
tests\test_compute_backend.py
tests\test_trajectory_hybrid_equivalence.py
tests\test_tracknet_hybrid.py
tests\test_startup_checks.py
tests\test_fusion_golden_telemetry.py
tests\test_wasb_runner.py
tests\test_tracknet_runner.py
tests\test_wsl_tilde_quoting.py
tests\test_stub_runner.py
tests\test_run_telemetry.py
tests\test_profile_parity.py
tests\test_output_base.py
tests\test_fusion_hybrid.py
tests\test_detector_keying.py
```

45. assistant (2026-07-09T06:32:35.450Z)

Let me read the two most relevant existing test files to match their mocking patterns.

46. user (2026-07-09T06:32:36.272Z)

```
1      """Worker dispatcher: pull (storage) ? run pipeline ? push (storage).
2
3      `cmd_infer` is exercised end-to-end over the LOCAL storage backend with the validator +
4      segmenter mocked, so it needs no real video/GPU/network. The real S3 round-trip is
5      validated
6      on the droplet against MinIO.
```

```

6     """
7
8     import json
9     import os
10
11     import pytest
12
13     from backend.config import load_config
14     from backend.results import Failure
15     from backend.worker import __main__ as wmain
16
17     # ----- pure helpers
18
19
20     @pytest.mark.parametrize(
21         "raw,expected",
22         [
23             ("true", True),
24             ("False", False),
25             ("3", 3),
26             ("2.5", 2.5),
27             ("wasb_hybrid", "wasb_hybrid"),
28         ],
29     )
30     def test_coerce(raw, expected):
31         assert wmain._coerce(raw) == expected and type(wmain._coerce(raw)) is type(expected)
32
33
34     def test_apply_overrides_on_typed_config():
35         cfg = _fresh_config()
36         wmain._apply_overrides(
37             cfg,
38             [
39                 "indexing.skip_ai_handoff=true",
40                 "indexing.min_segment_duration=1.5",
41                 "sport=tennis",
42             ],
43         )
44         assert cfg.indexing.skip_ai_handoff is True
45         assert cfg.indexing.min_segment_duration == 1.5
46         assert cfg.sport == "tennis"
47
48
49     def test_apply_overrides_rejects_bad_format():
50         cfg = _fresh_config()
51         with pytest.raises(ValueError):
52             wmain._apply_overrides(cfg, ["no_equals_sign"])
53
54
55     def test_write_intervals_csv(tmp_path):
56         segs = [
57             {
58                 "start_time": 2.0,
59                 "end_time": 5.0,
60                 "shot_count": 4,
61                 "shot_count_confidence": "LocalNoAI",
62             },
63             {"start_time": 9.0, "end_time": 12.5},
64         ]
65         out = tmp_path / "intervals.csv"
66         wmain._write_intervals_csv(segs, str(out))
67         lines = out.read_text().splitlines()
68         assert lines[0] == "start,end,shot_count,ending_reason"
69         assert lines[1].startswith("2.000,5.000,4,")
70         assert lines[2].startswith("9.000,12.500,")

```

```

71
72
73 def test_write_report_json(tmp_path):
74     out = tmp_path / "report.json"
75     wmain._write_report_json(
76         "vid1", [{"start_time": 1.0, "end_time": 2.0}], "success", str(out)
77     )
78     data = json.loads(out.read_text())
79     assert data["video_id"] == "vid1" and data["segment_count"] == 1
80     assert data["segments"][0] == {"start": 1.0, "end": 2.0}
81
82
83 def test_parser_infer_accepts_set_and_uris():
84     args = wmain.build_parser().parse_args(
85         [
86             "infer",
87             "--video-uri",
88             "s3://b/v.mp4",
89             "--out-uri",
90             "s3://b",
91             "--set",
92             "indexing.skip_ai_handoff=true",
93             "--set",
94             "sport=tennis",
95         ]
96     )
97     assert args.cmd == "infer" and args.video_uri == "s3://b/v.mp4"
98     assert args.set == ["indexing.skip_ai_handoff=true", "sport=tennis"]
99
100
101 def test_parser_infer_skip_validation_flag():
102     base = ["infer", "--video-uri", "s3://b/v.mp4", "--out-uri", "s3://b"]
103     # Default-OFF: the direct `worker infer` CLI still validates.
104     assert wmain.build_parser().parse_args(base).skip_validation is False
105     # The dispatcher passes it to disable the redundant worker re-validation.
106     assert wmain.build_parser().parse_args(base + ["--skip-validation"]).skip_validation
107     is True
108
109     # ----- cmd_infer (local,
110     mocked)
111
112     class _OKResult:
113         passed = True
114         validator_name = "fake"
115         message = "ok"
116
117         def model_dump(self):
118             return {"validator_name": "fake", "passed": True, "message": "ok"}
119
120
121     class _FailResult(_OKResult):
122         passed = False
123         message = "too blurry"
124
125
126     class _FakeSeg:
127         def __init__(self, db, config):
128             pass
129
130         def process_video(self, video_id, path, max_frames=None, progress=None):
131             # progress accepted for #534 live heartbeat compat (ignored in fake)
132             return [
133                 {

```

```

134         "start_time": 2.0,
135         "end_time": 5.0,
136         "shot_count": 4,
137         "shot_count_confidence": "LocalNoAI",
138     },
139     {
140         "start_time": 9.0,
141         "end_time": 12.0,
142         "shot_count": 6,
143         "shot_count_confidence": "LocalNoAI",
144     },
145 ]
146
147
148 def _fresh_config():
149     # Load defaults without touching any cwd config.json.
150     import tempfile
151
152     p = tempfile.NamedTemporaryFile("w", suffix=".json", delete=False)
153     p.write("{}")
154     p.close()
155     return load_config(p.name)
156
157
158 @pytest.fixture
159 def mocked_pipeline(monkeypatch):
160     monkeypatch.setattr(wmain, "compute_video_id", lambda p: "vid123")
161     monkeypatch.setattr(
162         wmain.validator_registry,
163         "run_all_with_context",
164         lambda path, cfg: ([_OKResult()], {}),
165     )
166     monkeypatch.setattr(wmain.segmenter_registry, "get", lambda name: _FakeSeg)
167     monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
168
169
170 def test_cmd_infer_local_roundtrip(tmp_path, mocked_pipeline):
171     video = tmp_path / "in.mp4"
172     video.write_bytes(b"FAKEVIDEO")
173     cfg = tmp_path / "config.json"
174     cfg.write_text("{}")
175     out = tmp_path / "bucket"
176     scratch = tmp_path / "scratch"
177
178     rc = wmain.main(
179         [
180             "infer",
181             "--video-uri",
182             str(video),
183             "--out-uri",
184             str(out),
185             "--config",
186             str(cfg),
187             "--scratch",
188             str(scratch),
189             "--segmenter",
190             "wasb_hybrid",
191             "--set",
192             "indexing.skip_ai_handoff=true",
193         ]
194     )
195     assert rc == 0
196     intervals = out / "outputs" / "vid123" / "intervals.csv"
197     report = out / "outputs" / "vid123" / "report.json"
198     assert intervals.exists() and report.exists()

```

```

199     rows = intervals.read_text().splitlines()
200     assert len(rows) == 3 # header + 2 rallies
201     assert json.loads(report.read_text())["segment_count"] == 2
202
203
204     class _EmptySeg:
205         """A segmenter that finds nothing ? a LEGITIMATE 0-rally match (a bare [])."""
206
207         def __init__(self, db, config):
208             pass
209
210         def process_video(self, video_id, path, max_frames=None, progress=None):
211             # progress accepted for #534 (ignored in fake)
212             return []
213
214
215     class _FailSeg:
216         """A segmenter whose DETECTOR failed/timed out ? returns a Failure (e.g. WASB
217 produced no
218 trajectory), which is NOT the same as a 0-rally match. This is the wf294 shape."""
219
220         def __init__(self, db, config):
221             pass
222
223         def process_video(self, video_id, path, max_frames=None, progress=None):
224             # progress accepted for #534 (ignored in fake)
225             return Failure(
226                 message="WASB detector produced no trajectory (timed out or aborted)",
227                 is_recoverable=True,
228             )
229
230     def test_setup_logging_levels():
231         import logging
232
233         wmain._setup_logging(False)
234         assert logging.getLogger().level == logging.INFO
235         wmain._setup_logging(True)
236         assert logging.getLogger().level == logging.DEBUG
237
238
239     def test_cmd_infer_zero_segments_is_logged_loudly(tmp_path, monkeypatch, caplog):
240         # The previous silent "0 segment(s)" gave nothing to analyse (the real cold-start
241 bug).
242         monkeypatch.setattr(wmain, "compute_video_id", lambda p: "vidZ")
243         monkeypatch.setattr(
244             wmain.validator_registry,
245             "run_all_with_context",
246             lambda path, cfg: ([_OKResult()], {}),
247         )
248         monkeypatch.setattr(wmain.segmenter_registry, "get", lambda name: _EmptySeg)
249         monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
250         monkeypatch.setenv("RALLY_DETECTOR_IMPL", "native")
251         video = tmp_path / "in.mp4"
252         video.write_bytes(b"FAKE")
253         out = tmp_path / "bucket"
254
255         # Call cmd_infer directly (not via main(), whose _setup_logging force-reconfigures
256 handlers and
257 # would detach caplog's). The warning is emitted by cmd_infer regardless of who
258 configures logging.
259         args = wmain.build_parser().parse_args(
260             [
261                 "infer",
262                 "--video-uri",

```

```

260         str(video),
261         "--out-uri",
262         str(out),
263         "--scratch",
264         str(tmp_path / "s"),
265         "--segmenter",
266         "wasb_hybrid",
267     ]
268 )
269 with caplog.at_level("WARNING", logger="backend.worker"):
270     rc = wmain.cmd_infer(args)
271     assert rc == 0 # still a clean run, just empty ? but now EXPLAINED
272     msgs = " ".join(r.getMessage() for r in caplog.records)
273     assert "0 segments" in msgs and "vidZ" in msgs
274     assert "detector_impl=native" in msgs # surfaces the resolved detector
275     assert "native runner UNAVAILABLE" in msgs # points at the likely cause
276     # A segmenter that returns a bare [] is a LEGITIMATE 0-rally match ? status stays
277     "success"
278     # (contrast test_cmd_infer_detector_failure_*, where a Failure surfaces as
279     status="failed").
280     report = json.loads((out / "outputs" / "vidZ" / "report.json").read_text())
281     assert report["segment_count"] == 0 and report["status"] == "success"
282
283 def test_cmd_infer_detector_failure_marks_report_and_marker_failed(
284     tmp_path, monkeypatch
285 ):
286     """wf294 fix: a segmenter that returns a Failure (detector timeout/abort) must NOT
287     be written
288     as a 0-segment "success". cmd_infer returns rc 3 and writes report.json + the done-
289     marker with
290     status="failed" so the control plane / alpha-user surface can show an error /
291     retry."""
292     monkeypatch.setattr(wmain, "compute_video_id", lambda p: "vidX")
293     monkeypatch.setattr(
294         wmain.validator_registry,
295         "run_all_with_context",
296         lambda path, cfg: ([_OKResult()], {}),
297     )
298     monkeypatch.setattr(wmain.segmenter_registry, "get", lambda name: _FailSeg)
299     monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
300     video = tmp_path / "in.mp4"
301     video.write_bytes(b"FAKE")
302     out = tmp_path / "bucket"
303     done = tmp_path / "bucket" / "_dispatch" / "tok.done"
304
305     rc = wmain.main(
306         [
307             "infer",
308             "--video-uri",
309             str(video),
310             "--out-uri",
311             str(out),
312             "--scratch",
313             str(tmp_path / "s"),
314             "--segmenter",
315             "wasb_hybrid",
316             "--done-uri",
317             str(done),
318         ]
319     )
320     # A detector failure is rc 3 (distinct from validation's rc 2 and a clean rc 0).
321     assert rc == 3
322     # report.json is written as an explicit FAILURE, not a false empty "success".
323     report = json.loads((out / "outputs" / "vidX" / "report.json").read_text())

```

```

320     assert report["status"] == "failed" and report["segment_count"] == 0
321     assert report["segments"] == []
322     # ...and the completion marker the cloud dispatcher bucket-polls carries the same
verdict.
323     marker = json.loads(done.read_text())
324     assert marker["returncode"] == 3 and marker["status"] == "failed"
325     assert marker["segment_count"] == 0 and marker["video_id"] == "vidX"
326
327
328     def test_cmd_infer_real_trajectory_segmenter_failure_flows_to_report(tmp_path,
monkeypatch):
329         """End-to-end: the ACTUAL WasbHybridSegmenter no-trajectory abort (the wf294 shape)
flows
330         through cmd_infer -> _fail(3) -> a status="failed" report.json + done-marker. Locks
the two
331         halves (segmenter Failure + worker report/marker) TOGETHER, not just via the
_FailSeg stub."""
332         from unittest.mock import MagicMock
333
334         from backend.pipeline.segmenters.wasb_hybrid import WasbHybridSegmenter
335
336         runner = MagicMock()
337         runner.healthcheck.return_value = (True, "env ok")
338         runner.run_predict.return_value = None # WASB timed out / aborted: no trajectory
produced
339
340         # Force the real base process_video down to its no-trajectory abort without a
GPU/video.
341         monkeypatch.setattr(WasbHybridSegmenter, "_resolve_runner", lambda self, cfg:
(runner, {}))
342         monkeypatch.setattr(
343             WasbHybridSegmenter, "_probe_fps_width", lambda *a, **k: (30.0, 1280.0)
344         )
345         monkeypatch.setattr(
346             "backend.pipeline.segmenters.trajectory_hybrid.get_video_duration", lambda p:
30.0
347         )
348         monkeypatch.setattr(wmain, "compute_video_id", lambda p: "vidR")
349         monkeypatch.setattr(
350             wmain.validator_registry, "run_all_with_context", lambda path, cfg:
([_OKResult()], {}))
351         )
352         monkeypatch.setattr(wmain.segmenter_registry, "get", lambda name:
WasbHybridSegmenter)
353         monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
354
355         video = tmp_path / "in.mp4"
356         video.write_bytes(b"FAKE")
357         out = tmp_path / "bucket"
358         done = tmp_path / "bucket" / "_dispatch" / "tok.done"
359
360         rc = wmain.main(
361             [
362                 "infer",
363                 "--video-uri",
364                 str(video),
365                 "--out-uri",
366                 str(out),
367                 "--scratch",
368                 str(tmp_path / "s"),
369                 "--segmenter",
370                 "wasb_hybrid",
371                 "--done-uri",
372                 str(done),
373                 "--set",

```

```

374         "indexing.skip_ai_handoff=true", # no API key needed; isolate the detector
abort
375     ]
376 )
377 assert rc == 3
378 report = json.loads((out / "outputs" / "vidR" / "report.json").read_text())
379 assert report["status"] == "failed" and report["segment_count"] == 0
380 marker = json.loads(done.read_text())
381 assert marker["returncode"] == 3 and marker["status"] == "failed"
382
383
384 def test_cmd_infer_validation_failure_returns_2(tmp_path, monkeypatch):
385     monkeypatch.setattr(wmain, "compute_video_id", lambda p: "vidF")
386     monkeypatch.setattr(
387         wmain.validator_registry,
388         "run_all_with_context",
389         lambda path, cfg: ([_FailResult()], {}),
390     )
391     monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
392     video = tmp_path / "in.mp4"
393     video.write_bytes(b"X")
394     cfg = tmp_path / "config.json"
395     cfg.write_text("{}")
396     out = tmp_path / "o"
397     done = tmp_path / "o" / "_dispatch" / "tok.done"
398     rc = wmain.main(
399         [
400             "infer",
401             "--video-uri",
402             str(video),
403             "--out-uri",
404             str(out),
405             "--config",
406             str(cfg),
407             "--scratch",
408             str(tmp_path / "s"),
409             "--done-uri",
410             str(done),
411         ]
412     )
413     assert rc == 2
414     # A validation failure is ALSO written as an explicit failure (not a missing
artifact / false
415     # success): report.json status="failed" + the done-marker rc 2 / status="failed".
416     report = json.loads((out / "outputs" / "vidF" / "report.json").read_text())
417     assert report["status"] == "failed" and report["segment_count"] == 0
418     marker = json.loads(done.read_text())
419     assert marker["returncode"] == 2 and marker["status"] == "failed"
420
421
422 def test_cmd_infer_skip_validation_bypasses_validators(tmp_path, monkeypatch):
423     """--skip-validation (set by the dispatcher, the authoritative gate) must NOT run
the worker's
424     validators ? no second decode, and no chance of a threshold mismatch REJECTING a
clip the
425     control-plane already passed (the live 2026-07-07 CUJ failure). Prove it by making
426     run_all_with_context RAISE: a skipped run never touches it, so the job completes rc
0 with a
427     status="success" report. Contrast test_cmd_infer_validation_failure_returns_2 (no
flag ? rc 2)."""
428     monkeypatch.setattr(wmain, "compute_video_id", lambda p: "vidSkip")
429
430     def _boom(*a, **k): # the validators must never be reached on the skip path
431         raise AssertionError("run_all_with_context must not run under --skip-
validation")

```



```

432
433 monkeypatch.setattr(wmain.validator_registry, "run_all_with_context", _boom)
434 monkeypatch.setattr(wmain.segmenter_registry, "get", lambda name: _FakeSeg)
435 monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
436 video = tmp_path / "in.mp4"
437 video.write_bytes(b"FAKE")
438 out = tmp_path / "bucket"
439
440 rc = wmain.main(
441     [
442         "infer",
443         "--video-uri",
444         str(video),
445         "--out-uri",
446         str(out),
447         "--scratch",
448         str(tmp_path / "s"),
449         "--segmenter",
450         "wasb_hybrid",
451         "--skip-validation",
452     ]
453 )
454 assert rc == 0 # a clip that WOULD fail re-validation still succeeds ? the gate is
trusted
455 report = json.loads((out / "outputs" / "vidSkip" / "report.json").read_text())
456 assert report["status"] == "success" and report["segment_count"] == 2
457
458
459 # --- cmd_train: pull labels -> synth trajectories -> manifest -> harness -> push bundle
---
460
461
462 @pytest.mark.parametrize(
463     "fname,expected",
464     [
465         (
466             "2026-06-22_GX020094.rallies.csv",
467             "GX020094",
468         ), # canonical dated label (DATA_IN_GCS §7.1)
469         ("GX020094.rallies.csv", "GX020094"), # legacy bare label -> SAME clip stem
470         (
471             "GX020094_proxy.rallies.csv",
472             "GX020094_proxy",
473         ), # _proxy is part of the name, kept
474         (
475             "2026-06-21_mahadevpura_1.rallies.csv",
476             "mahadevpura_1",
477         ), # underscores in <name> survive
478         ("v1.csv", "v1"), # bare .csv suffix
479         ("2026-06-22_GX020094.csv", "GX020094"), # dated + bare .csv
480     ],
481 )
482 def test_label_clip_name_strips_date_prefix(fname, expected):
483     """A dated label and a bare label must both resolve to the bare videos/<name>
stem."""
484     assert wmain._label_clip_name(fname) == expected
485
486
487 def test_cmd_train_detector_joins_dated_labels_to_bare_videos(tmp_path, monkeypatch):
488     """Regression (DATA_IN_GCS §7b blocker #1): canonical labels carry an ISO-date
prefix
489     (labels/<date>_<name>.rallies.csv) while videos are bare (videos/<name>.mp4). The
label-derived
490     clip name must strip the date so the video join succeeds ? otherwise every clip is
skipped and

```

```

491     train aborts with <2 videos."""
492     import os
493     import subprocess
494
495     monkeypatch.setenv("RALLY_DETECTOR_IMPL", "stub") # CPU stub runner (no GPU/WSL)
496     monkeypatch.setattr(wmain, "_probe_video_fps_width", lambda p: (30.0, 1280.0))
497     corpus = tmp_path / "corpus"
498     labels_dir = corpus / "labels"
499     videos_dir = corpus / "videos"
500     labels_dir.mkdir(parents=True)
501     videos_dir.mkdir(parents=True)
502     header = "rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
503     # canonical DATED labels ...
504     (labels_dir / "2026-06-22_GX020094.rallies.csv").write_text(
505         header + "1,2.0,5.0,winner,badminton,\n"
506     )
507     (labels_dir / "2026-06-21_mahadevpura_1.rallies.csv").write_text(
508         header + "1,3.0,6.0,winner,badminton,\n"
509     )
510     # ... joining to BARE-stem videos
511     (videos_dir / "GX020094.mp4").write_bytes(b"FAKEVIDEO1")
512     (videos_dir / "mahadevpura_1.mp4").write_bytes(b"FAKEVIDEO2")
513     out_dir = tmp_path / "out"
514     scratch = tmp_path / "scratch"
515     cfg = tmp_path / "config.json"
516     cfg.write_text("{}")
517
518     def fake_run(cmd, *a, **k):
519         out = cmd[cmd.index("--out") + 1]
520         ver = cmd[cmd.index("--version") + 1]
521         bdir = os.path.join(out, ver)
522         os.makedirs(bdir, exist_ok=True)
523         for fn in ("weights.json", "manifest.json"):
524             with open(os.path.join(bdir, fn), "w") as f:
525                 f.write("{}")
526         return subprocess.CompletedProcess(cmd, 0)
527
528     monkeypatch.setattr(subprocess, "run", fake_run)
529
530     rc = wmain.main(
531         [
532             "train",
533             "--corpus-uri",
534             str(corpus),
535             "--out-uri",
536             str(out_dir),
537             "--version",
538             "0.0.4",
539             "--config",
540             str(cfg),
541             "--scratch",
542             str(scratch),
543             "--trajectory-source",
544             "detector",
545             "--segmenter",
546             "wasb_hybrid",
547         ]
548     )
549     assert rc == 0 # before the fix: every clip skipped -> rc 2 ("<2 videos")
550     specs = json.loads((scratch / "manifest.json").read_text())
551     assert len(specs) == 2 # both dated labels joined
552     assert {s["name"] for s in specs} == {
553         "GX020094",
554         "mahadevpura_1",
555     } # downstream name is date-stripped

```

```

556
557
558     def test_cmd_train_detector_skips_label_with_no_matching_video(tmp_path, monkeypatch):
559         """A dated label whose <name> has no videos/<name>.<ext> is SKIPPED (the join-miss
branch),
560         while the labels that DO join are kept ? so one orphan label can't abort an
otherwise-valid run."""
561         import os
562         import subprocess
563
564         monkeypatch.setenv("RALLY_DETECTOR_IMPL", "stub")
565         monkeypatch.setattr(wmain, "_probe_video_fps_width", lambda p: (30.0, 1280.0))
566         corpus = tmp_path / "corpus"
567         labels_dir = corpus / "labels"
568         videos_dir = corpus / "videos"
569         labels_dir.mkdir(parents=True)
570         videos_dir.mkdir(parents=True)
571         header = "rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
572         row = "1,2.0,5.0,winner,badminton,\n"
573         (labels_dir / "2026-06-22_GX020094.rallies.csv").write_text(header + row)
574         (labels_dir / "2026-06-21_mahadevpura_1.rallies.csv").write_text(header + row)
575         (labels_dir / "2026-06-20_orphan.rallies.csv").write_text(
576             header + row
577         ) # NO videos/orphan.*
578         (videos_dir / "GX020094.mp4").write_bytes(b"FAKEVIDE01")
579         (videos_dir / "mahadevpura_1.mp4").write_bytes(b"FAKEVIDE02")
580         out_dir = tmp_path / "out"
581         scratch = tmp_path / "scratch"
582         cfg = tmp_path / "config.json"
583         cfg.write_text("{}")
584
585         def fake_run(cmd, *a, **k):
586             out = cmd[cmd.index("--out") + 1]
587             ver = cmd[cmd.index("--version") + 1]
588             bdir = os.path.join(out, ver)
589             os.makedirs(bdir, exist_ok=True)
590             for fn in ("weights.json", "manifest.json"):
591                 with open(os.path.join(bdir, fn), "w") as f:
592                     f.write("{}")
593             return subprocess.CompletedProcess(cmd, 0)
594
595         monkeypatch.setattr(subprocess, "run", fake_run)
596
597         rc = wmain.main(
598             [
599                 "train",
600                 "--corpus-uri",
601                 str(corpus),
602                 "--out-uri",
603                 str(out_dir),
604                 "--version",
605                 "0.0.5",
606                 "--config",
607                 str(cfg),
608                 "--scratch",
609                 str(scratch),
610                 "--trajectory-source",
611                 "detector",
612                 "--segmenter",
613                 "wasb_hybrid",
614             ]
615         )
616         assert (
617             rc == 0
618         ) # the orphan label was skipped; the two joinable clips carried the run

```

```

619     specs = json.loads((scratch / "manifest.json").read_text())
620     assert {s["name"] for s in specs} == {
621         "GX020094",
622         "mahadevpura_1",
623     } # 'orphan' absent
624
625
626     def test_cmd_train_detector_skips_unprobeable_and_trajless_videos(
627         tmp_path, monkeypatch
628     ):
629         """cmd_train's detector loop must SKIP (never crash, never guess) a video whose
630         fps/width won't
631         probe and a video the detector yields no trajectory for, keeping the healthy
632         clips."""
633
634         import os
635         import subprocess
636
637         corpus = tmp_path / "corpus"
638         labels_dir = corpus / "labels"
639         videos_dir = corpus / "videos"
640         labels_dir.mkdir(parents=True)
641         videos_dir.mkdir(parents=True)
642         header = "rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
643         row = "1,2.0,5.0,winner,badminton,\n"
644         for stem in ("good1", "good2", "badfps", "notraj"):
645             (labels_dir / f"2026-06-22_{stem}.rallies.csv").write_text(header + row)
646             (videos_dir / f"{stem}.mp4").write_bytes(b"FAKE-" + stem.encode())
647
648         # probe fails ONLY for badfps -> that clip is skipped (a wrong fps mis-aligns every
649         label).
650
651         monkeypatch.setattr(
652             wmain,
653             "_probe_video_fps_width",
654             lambda p: None if "badfps" in os.path.basename(p) else (30.0, 1280.0),
655         )
656         monkeypatch.setattr(
657             wmain, "compute_video_id", lambda p: "id-" + os.path.basename(p)
658         )
659
660         class _FakeRunner:
661             # writes a real trajectory for everything EXCEPT notraj (empty result ->
662             skipped).
663
664             def run_predict(self, video, out_dir, video_id=None):
665                 if "notraj" in os.path.basename(video):
666                     return ""
667                 traj = os.path.join(out_dir, f"{video_id}_traj.csv")
668                 with open(traj, "w") as f:
669                     f.write("frame,x,y\n0,1,2\n")
670                 return traj
671
672         monkeypatch.setattr(
673             wmain, "_resolve_train_detector", lambda args, config: _FakeRunner()
674         )
675
676         out_dir = tmp_path / "out"
677         scratch = tmp_path / "scratch"
678         cfg = tmp_path / "config.json"
679         cfg.write_text("{}")
680
681         def fake_run(cmd, *a, **k):
682             out = cmd[cmd.index("--out") + 1]
683             ver = cmd[cmd.index("--version") + 1]
684             bdir = os.path.join(out, ver)
685             os.makedirs(bdir, exist_ok=True)
686             for fn in ("weights.json", "manifest.json"):

```

```

680         with open(os.path.join(bdir, fn), "w") as f:
681             f.write("{}")
682         return subprocess.CompletedProcess(cmd, 0)
683
684 monkeypatch.setattr(subprocess, "run", fake_run)
685
686 rc = wmain.main(
687     [
688         "train",
689         "--corpus-uri",
690         str(corpus),
691         "--out-uri",
692         str(out_dir),
693         "--version",
694         "0.0.6",
695         "--config",
696         str(cfg),
697         "--scratch",
698         str(scratch),
699         "--trajectory-source",
700         "detector",
701         "--segmenter",
702         "wasb_hybrid",
703     ]
704 )
705 assert rc == 0
706 specs = json.loads((scratch / "manifest.json").read_text())
707 assert {s["name"] for s in specs} == {"good1", "good2"} # badfps + notraj skipped
708
709
710 def test_cmd_train_detector_aborts_clean_when_fewer_than_two_videos_join(
711     tmp_path, monkeypatch
712 ):
713     """When labels join to <2 videos, train aborts cleanly with rc 2 (leave-one-video-
714 out needs >=2).
715     This is the guard the date-prefix bug used to trip for EVERY clip (DATA_IN_GCS $7b
716 #1); here it
717     fires legitimately because only one of two labels has a matching video."""
718     import subprocess
719
720     monkeypatch.setenv("RALLY_DETECTOR_IMPL", "stub")
721     monkeypatch.setattr(wmain, "_probe_video_fps_width", lambda p: (30.0, 1280.0))
722     corpus = tmp_path / "corpus"
723     labels_dir = corpus / "labels"
724     videos_dir = corpus / "videos"
725     labels_dir.mkdir(parents=True)
726     videos_dir.mkdir(parents=True)
727     header = "rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
728     row = "1,2.0,5.0,winner,badminton,\n"
729     (labels_dir / "2026-06-22_GX020094.rallies.csv").write_text(header + row) # joins
730     (labels_dir / "2026-06-21_orphan.rallies.csv").write_text(
731         header + row
732     ) # no video -> skipped
733     (videos_dir / "GX020094.mp4").write_bytes(b"FAKEVIDE01")
734     cfg = tmp_path / "config.json"
735     cfg.write_text("{}")
736     # subprocess.run must NOT be reached (we abort before the harness); guard that.
737     monkeypatch.setattr(
738         subprocess, "run", lambda *a, **k: pytest.fail("harness ran despite <2 videos")
739     )
740
741 rc = wmain.main(
742     [
743         "train",
744         "--corpus-uri",

```

```

743         str(corpus),
744         "--out-uri",
745         str(tmp_path / "out"),
746         "--version",
747         "0.0.7",
748         "--config",
749         str(cfg),
750         "--scratch",
751         str(tmp_path / "s"),
752         "--trajectory-source",
753         "detector",
754         "--segmenter",
755         "wasb_hybrid",
756     ]
757 )
758 assert rc == 2 # only 1 of 2 labels joined -> below the leave-one-video-out floor
759
760
761 def test_cmd_train_orchestration(tmp_path, monkeypatch):
762     import os
763     import subprocess
764
765     labels_dir = tmp_path / "corpus" / "labels"
766     labels_dir.mkdir(parents=True)
767     header = "rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
768     (labels_dir / "v1.rallies.csv").write_text(
769         header + "1,2.0,5.0,winner,badminton,\n2,9.0,12.0,let,badminton,\n"
770     )
771     (labels_dir / "v2.rallies.csv").write_text(header + "1,3.0,6.0,winner,badminton,\n")
772     out_dir = tmp_path / "out"
773     scratch = tmp_path / "scratch"
774     cfg = tmp_path / "config.json"
775     cfg.write_text("{}")
776
777     # Mock the harness subprocess: write a fake versioned bundle at --out/--version.
778     def fake_run(cmd, *a, **k):
779         out = cmd[cmd.index("--out") + 1]
780         ver = cmd[cmd.index("--version") + 1]
781         bdir = os.path.join(out, ver)
782         os.makedirs(bdir, exist_ok=True)
783         for fn in ("weights.json", "manifest.json"):
784             with open(os.path.join(bdir, fn), "w") as f:
785                 f.write("{}")
786         return subprocess.CompletedProcess(cmd, 0)
787
788     monkeypatch.setattr(subprocess, "run", fake_run)
789
790     rc = wmain.main(
791         [
792             "train",
793             "--corpus-uri",
794             str(tmp_path / "corpus"),
795             "--out-uri",
796             str(out_dir),
797             "--version",
798             "0.0.1",
799             "--config",
800             str(cfg),
801             "--scratch",
802             str(scratch),
803         ]
804     )
805     assert rc == 0
806     # bundle pushed to <out-uri>/weights/<version>/
807     assert (out_dir / "weights" / "0.0.1" / "weights.json").exists()

```

```

808     assert (out_dir / "weights" / "0.0.1" / "manifest.json").exists()
809     # the manifest fed to the harness enumerated both labeled videos with a trajectory
each
810     specs = json.loads((scratch / "manifest.json").read_text())
811     assert len(specs) == 2
812     assert all("trajectory" in s and "labels" in s and "fps" in s for s in specs)
813
814
815     def test_parser_train_detector_args():
816         args = wmain.build_parser().parse_args(
817             [
818                 "train",
819                 "--corpus-uri",
820                 "gcs://b",
821                 "--out-uri",
822                 "gcs://b",
823                 "--version",
824                 "1.0.0",
825                 "--trajectory-source",
826                 "detector",
827                 "--segmenter",
828                 "fusion_hybrid",
829             ]
830         )
831         assert args.trajectory_source == "detector" and args.segmenter == "fusion_hybrid"
832
833
834     def test_cmd_train_detector_source_runs_real_extract_via_stub(tmp_path, monkeypatch):
835         """--trajectory-source detector pulls videos + runs the resolved DetectorRunner.
With
836         RALLY_DETECTOR_IMPL=stub the *wiring* is exercised GPU-free (the stub IS the
runner)."""
837         import os
838         import subprocess
839
840         monkeypatch.setenv(
841             "RALLY_DETECTOR_IMPL", "stub"
842         ) # resolve a CPU stub runner (no GPU/WSL)
843         # Fake bytes aren't decodable, so stub the (strict, fail-loud) ffprobe probe to a
real value.
844         monkeypatch.setattr(wmain, "_probe_video_fps_width", lambda p: (30.0, 1280.0))
845         corpus = tmp_path / "corpus"
846         labels_dir = corpus / "labels"
847         videos_dir = corpus / "videos"
848         labels_dir.mkdir(parents=True)
849         videos_dir.mkdir(parents=True)
850         header = "rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
851         (labels_dir / "v1.rallies.csv").write_text(header + "1,2.0,5.0,winner,badminton,\n")
852         (labels_dir / "v2.rallies.csv").write_text(header + "1,3.0,6.0,winner,badminton,\n")
853         (videos_dir / "v1.mp4").write_bytes(b"FAKEVIDE01")
854         (videos_dir / "v2.mp4").write_bytes(b"FAKEVIDE02")
855         (videos_dir / "v1.srt").write_text(
856             "1\n00:00:00,000 --> 00:00:01,000\nx\n"
857         ) # sidecar: must NOT shadow v1.mp4
858         out_dir = tmp_path / "out"
859         scratch = tmp_path / "scratch"
860         cfg = tmp_path / "config.json"
861         cfg.write_text("{}")
862
863         def fake_run(cmd, *a, **k):
864             out = cmd[cmd.index("--out") + 1]
865             ver = cmd[cmd.index("--version") + 1]
866             bdir = os.path.join(out, ver)
867             os.makedirs(bdir, exist_ok=True)
868             for fn in ("weights.json", "manifest.json"):

```

```

869         with open(os.path.join(bdir, fn), "w") as f:
870             f.write("{}")
871         return subprocess.CompletedProcess(cmd, 0)
872
873 monkeypatch.setattr(subprocess, "run", fake_run)
874
875 rc = wmain.main(
876     [
877         "train",
878         "--corpus-uri",
879         str(corpus),
880         "--out-uri",
881         str(out_dir),
882         "--version",
883         "0.0.2",
884         "--config",
885         str(cfg),
886         "--scratch",
887         str(scratch),
888         "--trajectory-source",
889         "detector",
890         "--segmenter",
891         "wasb_hybrid",
892     ]
893 )
894 assert rc == 0
895 specs = json.loads((scratch / "manifest.json").read_text())
896 assert (
897     len(specs) == 2
898     ) # the v1.srt sidecar did NOT shadow v1.mp4 (extension filter)
899 # Real-extract path: each spec points at an on-disk trajectory the runner actually
wrote.
900 assert all(os.path.exists(s["trajectory"]) for s in specs)
901 assert all(
902     s["trajectory"].endswith("_stub.csv") for s in specs
903     ) # the stub runner's output
904
905
906 def test_cmd_train_detector_native_on_tracknet_fails_clean_not_traceback(
907     tmp_path, monkeypatch
908 ):
909     """detector_impl=native for a family with no native runner must return rc 2 (clean),
not
910 propagate a raw NotImplementedError traceback out of the worker."""
911 monkeypatch.setenv("RALLY_DETECTOR_IMPL", "native")
912 corpus = tmp_path / "corpus"
913 labels_dir = corpus / "labels"
914 (labels_dir).mkdir(parents=True)
915 (corpus / "videos").mkdir(parents=True)
916 header = "rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
917 (labels_dir / "v1.rallies.csv").write_text(header + "1,2.0,5.0,winner,badminton,\n")
918 (labels_dir / "v2.rallies.csv").write_text(header + "1,3.0,6.0,winner,badminton,\n")
919 cfg = tmp_path / "config.json"
920 cfg.write_text("{}")
921 rc = wmain.main(
922     [
923         "train",
924         "--corpus-uri",
925         str(corpus),
926         "--out-uri",
927         str(tmp_path / "out"),
928         "--version",
929         "0.0.3",
930         "--config",
931         str(cfg),

```



```

932         "--scratch",
933         str(tmp_path / "s"),
934         "--trajectory-source",
935         "detector",
936         "--segmenter",
937         "tracknet_hybrid",
938     ]
939 )
940 assert (
941     rc == 2
942 ) # graceful: _resolve_train_detector caught NotImplementedError, returned None
943
944
945 def test_probe_video_fps_width_returns_none_on_unreadable(tmp_path):
946     """The strict ffprobe probe must return None (caller skips) for a non-video file,
947     never guess."""
948     bad = tmp_path / "not_a_video.mp4"
949     bad.write_bytes(b"NOTAVIDEO")
950     assert wmain._probe_video_fps_width(str(bad)) is None
951
952     # ----- compile subcommand (#521 ?
953 L4 stitch)
954
955 def test_parser_compile_accepts_spec_and_done_uris():
956     args = wmain.build_parser().parse_args(
957         [
958             "compile",
959             "--spec-uri",
960             "gs://b/_compile/tok/spec.json",
961             "--out-uri",
962             "gs://b",
963             "--done-uri",
964             "gs://b/_compile/tok.done",
965             "--set",
966             "storage.output_root=gs://b",
967         ]
968     )
969     assert args.cmd == "compile"
970     assert args.spec_uri == "gs://b/_compile/tok/spec.json"
971     assert args.done_uri == "gs://b/_compile/tok.done"
972     assert args.set == ["storage.output_root=gs://b"]
973
974
975 def _write_compile_spec(
976     bucket, video_id, source, name, encoder="h264_nvenc", rendition="original"
977 ):
978     spec = {
979         "video_id": video_id,
980         "video_uri": str(source),
981         "output_name": name,
982         "rendition": rendition,
983         "locale": None,
984         "time_pairs": [[2.0, 6.0], [10.0, 14.0]],
985         "stitching": {"encoder": encoder, "watermark": {"enabled": False}},
986     }
987     spec_dir = os.path.join(str(bucket), "_compile", "tok")
988     os.makedirs(spec_dir, exist_ok=True)
989     spec_uri = os.path.join(spec_dir, "spec.json")
990     with open(spec_uri, "w", encoding="utf-8") as f:
991         json.dump(spec, f)
992     return spec_uri
993
994

```

```

995     def test_cmd_compile_stitches_and_writes_marker(tmp_path, monkeypatch):
996         """cmd_compile reads the staged spec, runs the SHARED _stitch_reel core
(VideoCompiler stubbed so
997         no ffmpeg), mirrors the reel, and writes a compile done-marker for the dispatcher's
bucket-poll."""
998         import backend.orchestration as orch
999
1000         bucket = tmp_path / "bucket"
1001         bucket.mkdir()
1002         source = tmp_path / "src.mp4"
1003         source.write_bytes(b"SRC")
1004         spec_uri = _write_compile_spec(bucket, "vidReel", source, "reel.mp4")
1005         done = str(bucket / "_compile" / "tok.done")
1006
1007         seen = {}
1008
1009         def _fake_stitch(
1010             cfg,
1011             stitching,
1012             video_id,
1013             source_ref,
1014             time_pairs,
1015             out_name,
1016             locale,
1017             notify,
1018             rendition="original",
1019             require_mirror=False,
1020             **kw,
1021         ):
1022             seen["encoder"] = stitching.encoder
1023             seen["pairs"] = time_pairs
1024             seen["source"] = source_ref
1025             seen["rendition"] = rendition
1026             seen["require_mirror"] = require_mirror
1027             notify("stitching")
1028             return {
1029                 "status": "success",
1030                 "video_id": video_id,
1031                 "download_url": f"/api/highlights/{video_id}/{out_name}",
1032                 "reel_uri": f"gs://b/highlights/{video_id}/{out_name}",
1033                 "filename": out_name,
1034                 "rendition": rendition,
1035             }
1036
1037         monkeypatch.setattr(orch, "_stitch_reel", _fake_stitch)
1038
1039         cfg = tmp_path / "c.json"
1040         cfg.write_text('{"storage": {"backend": "local"}}')
1041         rc = wmain.cmd_compile(
1042             wmain.build_parser().parse_args(
1043                 [
1044                     "compile",
1045                     "--spec-uri",
1046                     spec_uri,
1047                     "--out-uri",
1048                     str(bucket),
1049                     "--done-uri",
1050                     done,
1051                     "--config",
1052                     str(cfg),
1053                     "--scratch",
1054                     str(tmp_path / "s"),
1055                     "--set",
1056                     f"storage.output_root={bucket}",
1057                 ]

```

```

1058         )
1059     )
1060     assert rc == 0
1061     # the shared stitch core got the resolved pairs + the NVENC encoder from the spec
1062     assert seen["encoder"] == "h264_nvenc"
1063     assert seen["pairs"] == [(2.0, 6.0), (10.0, 14.0)]
1064     assert seen["source"] == str(source)
1065     assert seen["require_mirror"] is True
1066     assert seen["rendition"] == "original"
1067     marker = json.loads(open(done, encoding="utf-8").read())
1068     assert marker["returncode"] == 0 and marker["status"] == "success"
1069     assert marker["download_url"] == "/api/highlights/vidReel/reel.mp4"
1070     assert marker["filename"] == "reel.mp4"
1071     assert marker["rendition"] == "original"
1072
1073
1074     def test_cmd_compile_sd_rendition_reaches_stitch_and_marker(tmp_path, monkeypatch):
1075         """The worker honors the staged SD rendition directly; it does not build an original
1076         reel first."""
1077         import backend.orchestration as orch
1078
1079         bucket = tmp_path / "bucket"
1080         bucket.mkdir()
1081         source = tmp_path / "src.mp4"
1082         source.write_bytes(b"SRC")
1083         spec_uri = _write_compile_spec(
1084             bucket, "vidReel", source, "reel-sd.mp4", rendition="sd"
1085         )
1086         done = str(bucket / "_compile" / "tok.done")
1087         seen = {}
1088
1089         def _fake_stitch(*args, **kwargs):
1090             seen["rendition"] = kwargs.get("rendition") or (args[-1] if args else
1091 "original")
1092             return {
1093                 "status": "success",
1094                 "video_id": "vidReel",
1095                 "download_url": "/api/highlights/vidReel/reel-sd.mp4",
1096                 "reel_uri": "gs://b/highlights/vidReel/reel-sd.mp4",
1097                 "filename": "reel-sd.mp4",
1098                 "rendition": args[-1],
1099             }
1100
1101         monkeypatch.setattr(orch, "_stitch_reel", _fake_stitch)
1102
1103         cfg = tmp_path / "c.json"
1104         cfg.write_text('{"storage": {"backend": "local"}}')
1105         rc = wmain.cmd_compile(
1106             wmain.build_parser().parse_args(
1107                 [
1108                     "compile",
1109                     "--spec-uri",
1110                     spec_uri,
1111                     "--out-uri",
1112                     str(bucket),
1113                     "--done-uri",
1114                     done,
1115                     "--config",
1116                     str(cfg),
1117                     "--scratch",
1118                     str(tmp_path / "s"),
1119                 ]
1120             )
1121         )

```

```

1121     assert rc == 0
1122     assert seen["rendition"] == "sd"
1123     marker = json.loads(open(done, encoding="utf-8").read())
1124     assert marker["filename"] == "reel-sd.mp4"
1125     assert marker["rendition"] == "sd"
1126
1127
1128 def test_cmd_compile_failure_returns_rc3_and_failed_marker(tmp_path, monkeypatch):
1129     """A stitch failure surfaces as rc 3 + a failed marker (rc != 0) so the dispatcher's
poll
1130     terminates and the SPA sees a clean failure."""
1131     import backend.orchestration as orch
1132
1133     bucket = tmp_path / "bucket"
1134     bucket.mkdir()
1135     source = tmp_path / "src.mp4"
1136     source.write_bytes(b"SRC")
1137     spec_uri = _write_compile_spec(bucket, "vidReel", source, "reel.mp4")
1138     done = str(bucket / "_compile" / "tok.done")
1139
1140     monkeypatch.setattr(
1141         orch,
1142         "_stitch_reel",
1143         lambda *a, **k: {"status": "failed", "message": "ffmpeg blew up", "video_id":
"vidReel"},
1144     )
1145     cfg = tmp_path / "c.json"
1146     cfg.write_text('{"storage": {"backend": "local"}}')
1147     rc = wmain.cmd_compile(
1148         wmain.build_parser().parse_args(
1149             [
1150                 "compile",
1151                 "--spec-uri",
1152                 spec_uri,
1153                 "--out-uri",
1154                 str(bucket),
1155                 "--done-uri",
1156                 done,
1157                 "--config",
1158                 str(cfg),
1159                 "--scratch",
1160                 str(tmp_path / "s"),
1161             ]
1162         )
1163     )
1164     assert rc == 3
1165     marker = json.loads(open(done, encoding="utf-8").read())
1166     assert marker["returncode"] == 3 and marker["status"] == "failed"
1167
1168
1169 def test_cmd_compile_no_done_uri_writes_no_marker(tmp_path, monkeypatch):
1170     """DEFAULT-OFF: a direct-CLI stitch with no --done-uri writes no marker."""
1171     import backend.orchestration as orch
1172
1173     bucket = tmp_path / "bucket"
1174     bucket.mkdir()
1175     source = tmp_path / "src.mp4"
1176     source.write_bytes(b"SRC")
1177     spec_uri = _write_compile_spec(bucket, "vidReel", source, "reel.mp4")
1178
1179     monkeypatch.setattr(
1180         orch, "_stitch_reel", lambda *a, **k: {"status": "success", "video_id":
"vidReel"}
1181     )
1182     cfg = tmp_path / "c.json"

```

```

1183     cfg.write_text('{"storage": {"backend": "local"}}')
1184     rc = wmain.cmd_compile(
1185         wmain.build_parser().parse_args(
1186             [
1187                 "compile",
1188                 "--spec-uri",
1189                 spec_uri,
1190                 "--out-uri",
1191                 str(bucket),
1192                 "--scratch",
1193                 str(tmp_path / "s"),
1194                 "--config",
1195                 str(cfg),
1196             ]
1197         )
1198     )
1199     assert rc == 0
1200     assert not (bucket / "_compile" / "tok.done").exists()
1201
1202
1203     def test_cmd_compile_fails_when_reel_upload_fails(tmp_path, monkeypatch):
1204         """Review finding (silent data loss): drive cmd_compile through the REAL
1205         _stitch_reel (VideoCompiler
1206         stubbed so no ffmpeg) with a failing reel mirror-upload. The worker has no local
1207         fallback, so it MUST
1208         return rc 3 + a FAILED marker ? never a 'success' marker for a reel the SPA can't
1209         download (404)."""
1210         import backend.main as bmain
1211
1212         bucket = tmp_path / "bucket"
1213         bucket.mkdir()
1214         source = tmp_path / "src.mp4"
1215         source.write_bytes(b"SRC")
1216         # output_root is a BUCKET URI so reel_object_uri resolves + the mirror is attempted
1217         (then fails)
1218         spec_uri = _write_compile_spec(bucket, "vidReel", source, "reel.mp4")
1219         done = str(bucket / "_compile" / "tok.done")
1220
1221         class _FakeCompiler:
1222             def __init__(self, output_dir, *a, **k):
1223                 self.output_dir = output_dir
1224                 os.makedirs(output_dir, exist_ok=True)
1225
1226             def compile_highlights(self, src, pairs, name):
1227                 p = os.path.join(self.output_dir, os.path.basename(name))
1228                 with open(p, "wb") as f:
1229                     f.write(b"REELBYTES")
1230                 return p
1231
1232         monkeypatch.setattr(bmain, "VideoCompiler", _FakeCompiler)
1233         monkeypatch.setattr(wmain.storage, "acquire_local_input", lambda ref, sc:
1234         str(source))
1235         monkeypatch.setattr(wmain.storage, "cleanup_acquired_local_input", lambda *a, **k:
1236         None)
1237
1238         def _put(src_path, uri, **k):
1239             # the reel mirror to the bucket fails; the tiny marker put (a different uri)
1240             still succeeds
1241             if uri.endswith("reel.mp4"):
1242                 raise RuntimeError("gcs 503 mid-upload")
1243             # marker + any other put: copy locally so the marker is written
1244             import shutil
1245
1246             os.makedirs(os.path.dirname(uri), exist_ok=True)
1247             shutil.copy(src_path, uri)

```

```

1241         return uri
1242
1243     monkeypatch.setattr(wmain.storage, "put", _put)
1244
1245     cfg = tmp_path / "c.json"
1246     cfg.write_text('{"storage": {"backend": "local"}}')
1247     rc = wmain.cmd_compile(
1248         wmain.build_parser().parse_args(
1249             [
1250                 "compile",
1251                 "--spec-uri",
1252                 spec_uri,
1253                 "--out-uri",
1254                 str(bucket),
1255                 "--done-uri",
1256                 done,
1257                 "--config",
1258                 str(cfg),
1259                 "--scratch",
1260                 str(tmp_path / "s"),
1261                 "--set",
1262                 "storage.output_root=gs://testbucket525",
1263             ]
1264         )
1265     )
1266     assert rc == 3 # NOT a false success
1267     marker = json.loads(open(done, encoding="utf-8").read())
1268     assert marker["returncode"] == 3 and marker["status"] == "failed"
1269     assert "upload" in (marker.get("message") or "").lower()
1270
1271
1272     def test_cmd_infer_respects_detector_progress_sample_config(tmp_path, monkeypatch):
1273         """#534: worker cmd_infer reads the configurable sampling value and the path
1274         remains stable when sampling is set via --set (no crash, config reaches the
1275         closure)."""
1276         class _NoopSeg:
1277             def __init__(self, db, config):
1278                 pass
1279             def process_video(self, video_id, path, max_frames=None, progress=None):
1280                 if progress:
1281                     progress("detector 50% (2500/5000)")
1282                 return [{"start_time": 0.0, "end_time": 1.0}]
1283
1284         monkeypatch.setattr(wmain.segmenter_registry, "get", lambda name: _NoopSeg)
1285         monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
1286         monkeypatch.setattr(
1287             wmain.validator_registry,
1288             "run_all_with_context",
1289             lambda p, c: ([], {}),
1290         )
1291
1292         # Track that storage.put is reachable (we don't assert count here to keep test
1293         # stable)
1294         called = []
1295         monkeypatch.setattr(wmain.storage, "put", lambda *a, **k: called.append(1))
1296
1297         video = tmp_path / "in.mp4"
1298         video.write_bytes(b"fake")
1299         out = tmp_path / "bucket"
1300
1301         # Explicit small sample value must be accepted by the run
1302         rc = wmain.main(
1303             [
1304                 "infer",
1305                 "--video-uri",

```

```

1304         str(video),
1305         "--out-uri",
1306         str(out),
1307         "--scratch",
1308         str(tmp_path / "s"),
1309         "--segmenter",
1310         "wasb_hybrid",
1311         "--skip-validation",
1312         "--set",
1313         "indexing.detector_progress_sample_sec=0.1",
1314     ]
1315 )
1316 assert rc == 0
1317 # At least the final serialize push happened (or the progress one)
1318 assert len(called) >= 0 # non-crashing is the main assertion here
1319
1320 # Large value also accepted
1321 rc = wmain.main(
1322     [
1323         "infer",
1324         "--video-uri",
1325         str(video),
1326         "--out-uri",
1327         str(out),
1328         "--scratch",
1329         str(tmp_path / "s2"),
1330         "--segmenter",
1331         "wasb_hybrid",
1332         "--skip-validation",
1333         "--set",
1334         "indexing.detector_progress_sample_sec=999",
1335     ]
1336 )
1337 assert rc == 0
1338

```

47. assistant (2026-07-09T06:33:10.595Z)

`test\_worker.py` gives me the exact mocking patterns (and confirms the fakes lack a `details` attr, so my score-hint must use `getattr(r, "details", {})`). Now let me read the native runner test to see how the #546 progress-streaming path is currently tested.

48. user (2026-07-09T06:33:11.986Z)

```

1      """NativeWasbRunner coverage ? the Linux-native (no-WSL) WASB detector path (M3.5).
2
3      The native subprocess is mocked (no torch/conda/GPU here). These tests lock the contract
4      the GPU-box byte-parity gate then relies on: it IS a DetectorRunner, it is NOT a WSL
runner
5      (no WslCommandMixin / wsl invocation), it builds a native argv (no `wsl`, no `conda
6      activate`, no `/mnt` translation) that runs the SAME wasb_infer.py with the SAME output
CSV
7      name as the WSL runner, threads the device through, and cleans the frame cache on
success.
8      """
9
10     import logging
11     import os
12     import types
13     from unittest.mock import patch
14
15     import pytest
16
17     from backend.pipeline.detectors.base import DetectorRunner, WslCommandMixin
18     from backend.pipeline.detectors.native_wasb_runner import (

```

```

19         NativeWasbConfig,
20         NativeWasbRunner,
21     )
22
23
24     def _proc(rc=0, stdout="", stderr=""):
25         return types.SimpleNamespace(returncode=rc, stdout=stdout, stderr=stderr)
26
27
28     # --- identity / no-WSL guarantee
29     -----
30     def test_is_a_detector_runner():
31         assert isinstance(NativeWasbRunner(NativeWasbConfig()), DetectorRunner)
32
33     def test_is_not_a_wsl_runner():
34         """The whole point of M3.5: the native runner must carry NO WSL plumbing."""
35         r = NativeWasbRunner(NativeWasbConfig())
36         assert not isinstance(r, WslCommandMixin)
37         assert not hasattr(r, "_wsl_bash")
38
39
40     def test_reuses_model_agnostic_windowing():
41         from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
42
43         assert NativeWasbRunner.parse_trajectory_csv is TrackNetRunner.parse_trajectory_csv
44         assert (
45             NativeWasbRunner.trajectory_to_action_windows
46             is TrackNetRunner.trajectory_to_action_windows
47         )
48
49
50     # --- config resolution (env overrides win over config over default)
51     -----
52     def test_from_indexing_cfg_defaults():
53         cfg = NativeWasbConfig.from_indexing_cfg({})
54         assert cfg.device == "cuda" and cfg.python_bin == "python"
55         # weights_path now comes from the Pydantic model default
56         (WasbIndexConfig.weights_path)
57         assert cfg.weights_path.endswith("wasb_badminton_best.pth.tar")
58         assert cfg.repo_dir == "~/models/WASB-SBDT"
59         # Native default is STREAMING (perf; bit-identical to disk, measured on a real GPU).
60         assert cfg.stream_video is True and cfg.keep_frames is False
61         # ?but an explicit stream_video=False is honoured (e.g. a container cv2 can't seek).
62         assert (
63             NativeWasbConfig.from_indexing_cfg(
64                 {"wasb": {"stream_video": False}}
65             ).stream_video
66             is False
67         )
68
69     def test_from_indexing_cfg_reads_config_block():
70         cfg = NativeWasbConfig.from_indexing_cfg(
71             {
72                 "wasb": {
73                     "weights_path": "/m/w.pth.tar",
74                     "device": "cpu",
75                     "python_bin": "/envs/wasb/bin/python",
76                     "repo_dir": "/m/WASB-SBDT",
77                     "stream_video": True,
78                     "keep_frames": True,
79                     "sport": "tennis",
80                 }
81             }

```



```

81         )
82         assert cfg.weights_path == "/m/w.pth.tar" and cfg.device == "cpu"
83         assert cfg.python_bin == "/envs/wasb/bin/python" and cfg.repo_dir == "/m/WASB-SBDT"
84         assert (
85             cfg.stream_video is True and cfg.keep_frames is True and cfg.sport == "tennis"
86         )
87
88
89     def test_env_overrides_beat_config(monkeypatch):
90         monkeypatch.setenv("WASB_WEIGHTS", "/env/w.pth.tar")
91         monkeypatch.setenv("WASB_DEVICE", "mps")
92         monkeypatch.setenv("WASB_PYTHON", "/env/py")
93         monkeypatch.setenv("WASB_REPO_DIR", "/env/repo")
94         monkeypatch.setenv("WASB_SCRATCH", "/env/scratch")
95         cfg = NativeWasbConfig.from_indexing_cfg(
96             {
97                 "wasb": {
98                     "weights_path": "/cfg/w",
99                     "device": "cpu",
100                     "python_bin": "/cfg/py",
101                     "repo_dir": "/cfg/repo",
102                 }
103             }
104         )
105         assert cfg.weights_path == "/env/w.pth.tar" and cfg.device == "mps"
106         assert cfg.python_bin == "/env/py" and cfg.repo_dir == "/env/repo"
107         assert cfg.scratch_dir == "/env/scratch"
108
109
110     # --- healthcheck
111     -----
112     def _ok_cfg(tmp_path, **kw):
113         w = tmp_path / "w.pth.tar"
114         w.write_text("weights")
115         (tmp_path / "repo" / "src").mkdir(parents=True)
116         return NativeWasbConfig(weights_path=str(w), repo_dir=str(tmp_path / "repo"), **kw)
117
118     def test_healthcheck_ok(tmp_path):
119         r = NativeWasbRunner(_ok_cfg(tmp_path, device="cuda"))
120         with patch.object(
121             r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda True")
122         ):
123             ok, msg = r.healthcheck()
124             assert ok and "torch" in msg
125
126
127     def test_healthcheck_empty_weights():
128         ok, msg = NativeWasbRunner(NativeWasbConfig(weights_path="")).healthcheck()
129         assert not ok and "weights" in msg.lower()
130
131
132     def test_healthcheck_weights_missing(tmp_path):
133         (tmp_path / "repo" / "src").mkdir(parents=True)
134         cfg = NativeWasbConfig(
135             weights_path=str(tmp_path / "nope.pth"), repo_dir=str(tmp_path / "repo")
136         )
137         ok, msg = NativeWasbRunner(cfg).healthcheck()
138         assert not ok and "not found" in msg.lower()
139
140
141     def test_healthcheck_repo_missing(tmp_path):
142         w = tmp_path / "w.pth.tar"
143         w.write_text("x")
144         cfg = NativeWasbConfig(weights_path=str(w), repo_dir=str(tmp_path / "absent"))

```

```

145         ok, msg = NativeWasbRunner(cfg).healthcheck()
146         assert not ok and "repo" in msg.lower()
147
148
149     def test_healthcheck_torch_import_fails(tmp_path):
150         r = NativeWasbRunner(_ok_cfg(tmp_path))
151         with patch.object(
152             r, "_run", return_value=_proc(1, stderr="ModuleNotFoundError: torch")
153         ):
154             ok, msg = r.healthcheck()
155             assert not ok and "failed" in msg.lower()
156
157
158     def test_healthcheck_cuda_unavailable_when_device_cuda(tmp_path):
159         r = NativeWasbRunner(_ok_cfg(tmp_path, device="cuda"))
160         with patch.object(
161             r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda False")
162         ):
163             ok, msg = r.healthcheck()
164             assert not ok and "cuda" in msg.lower()
165
166
167     def test_healthcheck_cpu_device_ok_without_cuda(tmp_path):
168         r = NativeWasbRunner(_ok_cfg(tmp_path, device="cpu"))
169         with patch.object(
170             r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda False")
171         ):
172             ok, msg = r.healthcheck()
173             assert ok # cpu run does not require a GPU
174
175
176     def test_healthcheck_python_not_found(tmp_path):
177         r = NativeWasbRunner(_ok_cfg(tmp_path, python_bin="/no/such/python"))
178         with patch.object(r, "_run", side_effect=FileNotFoundError()):
179             ok, msg = r.healthcheck()
180             assert not ok and "python" in msg.lower()
181
182
183     # --- M4: device="auto" CPU fallback (DeviceContext)
184     -----
185     def test_healthcheck_unknown_device_fails_fast(tmp_path):
186         """A typo'd device fails the healthcheck with a clear error, BEFORE the torch probe
187         subprocess
188         (beats a cryptic argparse rc=2 from wasb_infer.py at run time)."""
189         r = NativeWasbRunner(_ok_cfg(tmp_path, device="gpu"))
190         with patch.object(r, "_run") as spy:
191             ok, msg = r.healthcheck()
192             assert not ok and "unknown device" in msg.lower() and "gpu" in msg
193             spy.assert_not_called()
194
195
196     def test_healthcheck_auto_uses_cuda_when_available(tmp_path):
197         r = NativeWasbRunner(_ok_cfg(tmp_path, device="auto"))
198         with patch.object(
199             r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda True")
200         ):
201             ok, _msg = r.healthcheck()
202             assert ok
203             assert r._cuda_available is True
204             assert r._effective_device() == "cuda" # auto resolves to the GPU when present
205
206
207     def test_healthcheck_auto_falls_back_to_cpu_with_warning(tmp_path, caplog):
208         """device=auto on a GPU-less box does NOT hard-fail (unlike device=cuda) ? it
209         resolves to cpu

```

```

207         and warns loudly. This is the M4 portability win for a no-GPU Linux / cpu-serving
box. """
208         r = NativeWasbRunner(_ok_cfg(tmp_path, device="auto"))
209         with patch.object(
210             r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda False")
211         ), caplog.at_level(
212             logging.WARNING, logger="backend.pipeline.detectors.native_wasb_runner"
213         ):
214             ok, _msg = r.healthcheck()
215             assert ok # auto does not require a GPU
216             assert r._effective_device() == "cpu"
217             assert any("FALLING BACK TO CPU" in rec.message for rec in caplog.records)
218
219
220     def test_healthcheck_caches_probe_for_run_predict(tmp_path):
221         """healthcheck's CUDA probe is cached so run_predict reuses it (no second
subprocess). """
222         r = NativeWasbRunner(_ok_cfg(tmp_path, device="auto"))
223         with patch.object(
224             r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda False")
225         ) as spy:
226             r.healthcheck()
227             assert spy.call_count == 1
228             assert r._effective_device() == "cpu" # uses the cache, no extra probe
229             assert spy.call_count == 1
230
231
232     def test_run_predict_auto_threads_cpu_when_no_gpu(tmp_path):
233         """run_predict with device=auto + no GPU: the lazy probe (mocked empty stdout = no
cuda)
234         resolves to cpu, and 'auto' is NEVER sent to wasb_infer.py. """
235         _, _out, _rm, argv, _cwd, _key = _drive_success(
236             NativeWasbConfig(weights_path="/m/w.pth", device="auto"), tmp_path
237         )
238         assert argv[argv.index("--device") + 1] == "cpu"
239         assert "auto" not in argv
240
241
242     def test_run_predict_auto_threads_cuda_when_gpu(tmp_path):
243         """run_predict with device=auto + a GPU: the probe reports cuda True ? --device
cuda. """
244         r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth", device="auto"))
245         key = "clip__abc123abc123"
246         (tmp_path / f"{key}_wasb.csv").write_text("Frame,Visibility,X,Y\n")
247
248         def _run_side(argv, cwd=None):
249             if "-c" in argv: # the CUDA probe
250                 return _proc(0, stdout="cuda True")
251             return _proc(0) # the inference
252
253         with patch.object(r, "_copy_infer_script", return_value=True), patch.object(
254             r, "_run", side_effect=_run_side
255         ) as spy, patch.object(r, "_rm_rf"):
256             out = r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
257             assert out is not None
258             infer_argv = next(
259                 c.args[0] for c in spy.call_args_list if "wasb_infer.py" in c.args[0]
260             )
261             assert infer_argv[infer_argv.index("--device") + 1] == "cuda"
262
263
264     def test_run_predict_explicit_cuda_does_not_probe(tmp_path):
265         """An explicit device=cuda resolves to itself with NO probe subprocess ? unchanged
behaviour
266         (a single _run call: the inference). """

```

```

267     r, _out, _rm, argv, _cwd, _key = _drive_success(
268         NativeWasbConfig(weights_path="/m/w.pth", device="cuda"), tmp_path
269     )
270     assert argv[argv.index("--device") + 1] == "cuda"
271     assert r._cuda_available is None # never probed for an explicit device
272
273
274     # --- run_predict
275     -----
276     def _drive_success(cfg, tmp_path, vid="abc123abc123", **patches):
277         """Drive run_predict to success with the subprocess mocked; returns (runner, out,
278         rm_spy, argv, cwd)."""
279         r = NativeWasbRunner(cfg)
280         key = f"clip_{vid[:12]}"
281         (tmp_path / f"{key}_wasb.csv").write_text(
282             "Frame,Visibility,X,Y\n"
283         ) # success = CSV exists
284         with patch.object(r, "_copy_infer_script", return_value=True), patch.object(
285             r, "_run", return_value=_proc(0)
286         ) as run_spy, patch.object(r, "_rm_rf") as rm_spy:
287             out = r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id=vid)
288             argv = run_spy.call_args[0][0]
289             cwd = run_spy.call_args.kwargs.get("cwd")
290             return r, out, rm_spy, argv, cwd, key
291
292     def test_run_predict_builds_native_argv_no_wsl(tmp_path):
293         # Pin disk mode here (default is now streaming) so this also covers the
294         # --frames_out_dir path.
295         _, out, rm, argv, cwd, key = _drive_success(
296             NativeWasbConfig(weights_path="/m/w.pth", device="cuda", stream_video=False),
297             tmp_path,
298         )
299         assert out == os.path.join(os.path.abspath(str(tmp_path)), f"{key}_wasb.csv")
300         # No WSL / conda plumbing anywhere in the command.
301         assert "wsl" not in argv
302         assert not any("conda" in a or "/mnt/" in a for a in argv)
303         # It runs the SAME inference core, device-threaded, writing the SAME CSV name as
304         WSL.
305         assert argv[0] == "python" and argv[1] == "wasb_infer.py"
306         assert "--device" in argv and argv[argv.index("--device") + 1] == "cuda"
307         assert "--weights" in argv and argv[argv.index("--weights") + 1] == "/m/w.pth"
308         assert argv[argv.index("--out") + 1] == out
309         # Disk mode extracts frames + cleans them on success.
310         assert "--frames_out_dir" in argv and "--stream-video" not in argv
311         assert cwd.endswith(os.path.join("src"))
312         rm.assert_called_once()
313
314     def test_run_predict_default_is_stream_mode_no_frames_dir(tmp_path):
315         # No stream_video set -> the NATIVE DEFAULT (streaming): no PNG extraction, nothing
316         # to clean.
317         _, out, rm, argv, _cwd, _key = _drive_success(
318             NativeWasbConfig(weights_path="/m/w.pth"), tmp_path
319         )
320         assert out is not None
321         assert "--stream-video" in argv and "--frames_out_dir" not in argv
322         rm.assert_not_called() # streaming never materializes a frame cache
323
324     def test_run_predict_omits_gpu_feed_flags_by_default(tmp_path):
325         """batch_size/prefetch_batches default to 0 ? the flags are NOT passed, so
326         wasb_infer.py's
327         parity defaults (batch 8, no prefetch) and the byte-parity guarantee are
328         untouched."""

```

```

325     _, _out, _rm, argv, _cwd, _key = _drive_success(
326         NativeWasbConfig(weights_path="/m/w.pth"), tmp_path
327     )
328     assert "--batch-size" not in argv
329     assert "--prefetch-batches" not in argv
330
331
332     def test_run_predict_passes_log_every_batches_for_progress_sampling(tmp_path):
333         """#534: log_every_batches (emission cadence for detector %) is configurable and
334         threaded into the native argv so the sidecar sampling rate can be controlled end-to-
335         end."""
336         cfg = NativeWasbConfig(weights_path="/m/w.pth", log_every_batches=10)
337         _, _out, _rm, argv, _cwd, _key = _drive_success(cfg, tmp_path)
338         assert "--log-every-batches" in argv
339         assert argv[argv.index("--log-every-batches") + 1] == "10"
340
341         # With the new configurable default (50), the flag is always passed for native
342         # so the emission rate for progress/sidecar sampling is explicit and overridable.
343         cfg_default = NativeWasbConfig(weights_path="/m/w.pth")
344         _, _out, _rm, argv_def, _cwd, _key = _drive_success(cfg_default, tmp_path)
345         assert "--log-every-batches" in argv_def
346         assert argv_def[argv_def.index("--log-every-batches") + 1] == "50"
347
348     def test_run_predict_threads_gpu_feed_tuning_when_set(tmp_path):
349         """The cloud-serving preset's GPU-feed tuning reaches the wasb_infer.py argv."""
350         _, _out, _rm, argv, _cwd, _key = _drive_success(
351             NativeWasbConfig(
352                 weights_path="/m/w.pth", batch_size=64, prefetch_batches=4
353             ),
354             tmp_path,
355         )
356         assert argv[argv.index("--batch-size") + 1] == "64"
357         assert argv[argv.index("--prefetch-batches") + 1] == "4"
358
359
360     def test_run_predict_frames_dir_uses_scratch_and_is_cleaned(tmp_path):
361         """Frames extract under scratch_dir (the box's big NVMe) when set, and the SAME dir
362         is the one cleaned on success ? a regression that dropped scratch_dir (frames -> small
363         system disk) or diverged the cleanup target would be a disk-fill leak."""
364         scratch = tmp_path / "scratch"
365         r, out, rm, argv, _cwd, key = _drive_success(
366             NativeWasbConfig(
367                 weights_path="/m/w.pth", scratch_dir=str(scratch), stream_video=False
368             ),
369             tmp_path,
370         )
371         expected_frames = os.path.join(str(scratch), f"{key}_frames")
372         assert argv[argv.index("--frames_out_dir") + 1] == expected_frames
373         rm.assert_called_once_with(expected_frames) # cleanup target == extraction target
374
375
376     def test_run_predict_frames_dir_defaults_next_to_output(tmp_path):
377         _, out, rm, argv, _cwd, key = _drive_success(
378             NativeWasbConfig(weights_path="/m/w.pth", stream_video=False), tmp_path
379         )
380         expected_frames = os.path.join(os.path.abspath(str(tmp_path)), f"{key}_frames")
381         assert argv[argv.index("--frames_out_dir") + 1] == expected_frames
382         rm.assert_called_once_with(expected_frames)
383
384
385     def test_copy_infer_script_copies_byte_identical(tmp_path):
386         """Unmocked: the native runner stages the SAME wasb_infer.py the WSL runner runs

```

```

(the
387     byte-parity precondition) into <repo>/src/, byte-for-byte.""
388     import filecmp
389
390     from backend.pipeline.detectors import native_wasb_runner as nwr
391
392     r = NativeWasbRunner(NativeWasbConfig(repo_dir=str(tmp_path / "repo")))
393     assert r._copy_infer_script() is True
394     dest = tmp_path / "repo" / "src" / "wasb_infer.py"
395     assert dest.exists() and filecmp.cmp(nwr._INFER_PY, str(dest), shallow=False)
396
397
398     def test_copy_infer_script_false_on_oserror(tmp_path, monkeypatch):
399         r = NativeWasbRunner(NativeWasbConfig(repo_dir=str(tmp_path / "repo")))
400         monkeypatch.setattr(
401             "backend.pipeline.detectors.native_wasb_runner.shutil.copy",
402             lambda *a, **k: (_ for _ in ()).throw(OSError("no space")),
403         )
404         assert r._copy_infer_script() is False
405
406
407     def test_run_predict_device_threaded_cpu(tmp_path):
408         _, _out, _rm, argv, _cwd, _key = _drive_success(
409             NativeWasbConfig(weights_path="/m/w.pth", device="cpu"), tmp_path
410         )
411         assert argv[argv.index("--device") + 1] == "cpu"
412
413
414     def test_run_predict_keep_frames_skips_cleanup(tmp_path):
415         # Disk mode (stream_video=False) so the keep_frames branch is the reason cleanup is
416         # (under the default stream mode there are no frames to clean, which would be
417         # (vacuous).
418         _, _out, rm, _argv, _cwd, _key = _drive_success(
419             NativeWasbConfig(weights_path="/m/w.pth", keep_frames=True, stream_video=False),
420             tmp_path,
421         )
422         rm.assert_not_called()
423
424
425     def test_run_predict_none_when_copy_infer_fails(tmp_path):
426         r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth"))
427         with patch.object(r, "_copy_infer_script", return_value=False):
428             assert (
429                 r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
430                 is None
431             )
432
433
434     def test_run_predict_none_when_subprocess_fails(tmp_path):
435         r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth"))
436         with patch.object(r, "_copy_infer_script", return_value=True), patch.object(
437             r, "_run", return_value=_proc(1, stderr="CUDA error")
438         ):
439             assert (
440                 r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
441                 is None
442             )
443
444
445     def test_run_predict_none_when_csv_absent(tmp_path):
446         # Subprocess "succeeds" (rc 0) but writes no CSV -> treated as failure.
447         r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth"))
448         with patch.object(r, "_copy_infer_script", return_value=True), patch.object(
449             r, "_run", return_value=_proc(0)

```

```

449         ):
450             assert (
451                 r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
452                 is None
453             )
454
455
456     def test_run_predict_times_out_returns_none(tmp_path):
457         import subprocess
458
459         r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth"))
460         with patch.object(r, "_copy_infer_script", return_value=True), patch.object(
461             r, "_run", side_effect=subprocess.TimeoutExpired(cmd="x", timeout=1)
462         ):
463             assert (
464                 r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
465                 is None
466             )
467
468
469     def test_csv_name_matches_wsl_contract(tmp_path):
470         """The native CSV name must equal the WSL runner's so downstream + the parity diff
line up."""
471         _, out, _rm, _argv, _cwd, key = _drive_success(
472             NativeWasbConfig(weights_path="/m/w.pth"), tmp_path
473         )
474         assert os.path.basename(out) == f"{key}_wasb.csv"
475
476
477     # --- #506: GPU-resident NVDEC decode backend
-----
478     def test_run_predict_omits_decode_backend_flag_by_default(tmp_path):
479         """decode_backend defaults to "cpu" ? the flag is NOT passed, so the default argv
480         (and its byte-parity guarantee) is unchanged by #506."""
481         _, _out, _rm, argv, _cwd, _key = _drive_success(
482             NativeWasbConfig(weights_path="/m/w.pth"), tmp_path
483         )
484         assert "--decode-backend" not in argv
485
486
487     def test_run_predict_nvdec_resident_threads_flag_and_skips_cv2_paths(tmp_path):
488         """nvdec_resident decodes on the GPU inside wasb_infer.py: the flag is threaded and
489         NEITHER cv2 path flag (--stream-video / --frames_out_dir) rides along; there is no
490         frame cache to clean."""
491         _, out, rm, argv, _cwd, _key = _drive_success(
492             NativeWasbConfig(weights_path="/m/w.pth", decode_backend="nvdec_resident"),
493             tmp_path,
494         )
495         assert out is not None
496         assert argv[argv.index("--decode-backend") + 1] == "nvdec_resident"
497         assert "--stream-video" not in argv and "--frames_out_dir" not in argv
498         assert argv[argv.index("--device") + 1] == "cuda"
499         rm.assert_not_called()
500
501
502     def test_run_predict_nvdec_resident_auto_without_gpu_aborts_before_subprocess(tmp_path):
503         """device=auto resolving to cpu cannot host NVDEC decode: run_predict fails BEFORE
504         launching the (paid) inference subprocess ? only the CUDA probe runs."""
505         r = NativeWasbRunner(
506             NativeWasbConfig(
507                 weights_path="/m/w.pth", device="auto", decode_backend="nvdec_resident"
508             )
509         )
510         with patch.object(r, "_copy_infer_script", return_value=True), patch.object(
511             r, "_run", return_value=_proc(0, stdout="torch 2.6.0 cuda False")

```

```

512         ) as spy:
513             out = r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
514             assert out is None
515             assert spy.call_count == 1 # the probe only; wasb_infer.py never launched
516
517
518     def test_healthcheck_rejects_unknown_decode_backend(tmp_path):
519         """A typo'd decode_backend fails fast with a clear message, before any
520 subprocess."""
521         r = NativeWasbRunner(_ok_cfg(tmp_path, decode_backend="nvdec"))
522         with patch.object(r, "_run") as spy:
523             ok, msg = r.healthcheck()
524             assert not ok and "decode_backend" in msg
525             spy.assert_not_called()
526
527
528     def test_healthcheck_nvdec_resident_requires_effective_cuda(tmp_path):
529         """auto?cpu fallback is fine for the cpu decode path but NOT for nvdec_resident:
530 the healthcheck must hard-fail instead of warn-and-proceed."""
531         r = NativeWasbRunner(
532             _ok_cfg(tmp_path, device="auto", decode_backend="nvdec_resident")
533         )
534         with patch.object(
535             r, "_run", return_value=_proc(0, stdout="torch 2.6.0 cuda False")
536         ):
537             ok, msg = r.healthcheck()
538             assert not ok and "nvdec_resident" in msg
539
540
541     def test_healthcheck_nvdec_resident_ok_with_cuda_and_pynvc(tmp_path):
542         r = NativeWasbRunner(
543             _ok_cfg(tmp_path, device="cuda", decode_backend="nvdec_resident")
544         )
545         with patch.object(
546             r, "_run", return_value=_proc(0, stdout="torch 2.6.0 cuda True")
547         ) as spy:
548             ok, _msg = r.healthcheck()
549             assert ok
550             # Both probes ran: the torch/cuda one AND the PyNvVideoCodec import.
551             assert spy.call_count == 2
552             assert any("PyNvVideoCodec" in c.args[0][-1] for c in spy.call_args_list)
553
554
555     def test_healthcheck_nvdec_resident_requires_pynvc(tmp_path):
556         """A wasb env without PyNvVideoCodec must fail the healthcheck, not the paid
557 inference subprocess (the common misconfig: env built without the package)."""
558         r = NativeWasbRunner(
559             _ok_cfg(tmp_path, device="cuda", decode_backend="nvdec_resident")
560         )
561
562         def _run_side(argv, cwd=None):
563             if "PyNvVideoCodec" in argv[-1]:
564                 return _proc(1, stderr="ModuleNotFoundError: No module named
565 'PyNvVideoCodec'")
566             return _proc(0, stdout="torch 2.6.0 cuda True")
567
568         with patch.object(r, "_run", side_effect=_run_side):
569             ok, msg = r.healthcheck()
570             assert not ok and "PyNvVideoCodec" in msg
571
572
573     def test_run_predict_nvdec_resident_omits_prefetch_batches(tmp_path):
574         """prefetch overlaps CPU batch assembly with GPU compute ? inapplicable to the
575 GPU-resident decoder, so the dead knob must not ride into the argv (batch-size,
576 which DOES batch nvdec windows, still threads)."""

```



```

575     _, _out, _rm, argv, _cwd, _key = _drive_success(
576         NativeWasbConfig(
577             weights_path="/m/w.pth",
578             decode_backend="nvdec_resident",
579             batch_size=64,
580             prefetch_batches=4,
581         ),
582         tmp_path,
583     )
584     assert "--prefetch-batches" not in argv
585     assert argv[argv.index("--batch-size") + 1] == "64"
586
587
588     # --- the detector_impl='native' seam (parallels the stub seam)
589     -----
590     def test_resolve_runner_native_selects_native_runner_for_wasb(monkeypatch):
591         monkeypatch.delenv("RALLY_DETECTOR_IMPL", raising=False)
592         from backend.pipeline.segmenters.wasb_hybrid import WasbHybridSegmenter
593
594         runner, params = WasbHybridSegmenter(None, {}).resolve_runner(
595             {"detector_impl": "native"}
596         )
597         assert isinstance(runner, NativeWasbRunner)
598         assert params.get("detector_impl") == "native"
599
600     def test_resolve_runner_native_env_overrides_config(monkeypatch):
601         monkeypatch.setenv("RALLY_DETECTOR_IMPL", "native")
602         from backend.pipeline.segmenters.wasb_hybrid import WasbHybridSegmenter
603
604         runner, _ = WasbHybridSegmenter(None, {}).resolve_runner({"detector_impl": "auto"})
605         assert isinstance(runner, NativeWasbRunner)
606
607
608     def test_resolve_runner_native_supported_for_fusion_wasb(monkeypatch):
609         monkeypatch.delenv("RALLY_DETECTOR_IMPL", raising=False)
610         from backend.pipeline.segmenters.fusion_hybrid import FusionSegmenter
611
612         runner, params = FusionSegmenter(None, {}).resolve_runner(
613             {"detector_impl": "native"}
614         )
615         assert isinstance(runner, NativeWasbRunner)
616         assert params.get("fusion_substrate") == "wasb"
617
618
619     def test_resolve_runner_native_refused_for_fusion_tracknet(monkeypatch):
620         monkeypatch.delenv("RALLY_DETECTOR_IMPL", raising=False)
621         from backend.pipeline.segmenters.fusion_hybrid import FusionSegmenter
622
623         with pytest.raises(NotImplementedError):
624             FusionSegmenter(None, {}).resolve_runner(
625                 {"detector_impl": "native", "fusion": {"substrate": "tracknet"}}
626             )
627
628
629     def test_resolve_runner_native_refused_for_tracknet(monkeypatch):
630         monkeypatch.delenv("RALLY_DETECTOR_IMPL", raising=False)
631         from backend.pipeline.segmenters.tracknet_hybrid import TrackNetHybridSegmenter
632
633         with pytest.raises(NotImplementedError):
634             TrackNetHybridSegmenter(None, {}).resolve_runner({"detector_impl": "native"})
635
636
637     # --- progress callback for live detector heartbeat (#534)
638     -----

```

```

638 def test_run_predict_streams_and_calls_progress_on_pct_lines(tmp_path, caplog):
639     """When progress= is supplied, run_predict switches to Popen streaming and forwards
640     parsed "done/total (pct%)" lines (the format already emitted by wasb_infer) as
641     "detector XX% (D/T)" callbacks. Non-matching lines are still logged but do not call
progress.
642     The fast _run path is used when progress=None (byte-identical for tests)."""
643     import types
644
645     from backend.pipeline.detectors.native_wasb_runner import NativeWasbRunner
646
647     r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth"))
648     key = "clip__abc123abc123"
649     (tmp_path / f"{key}_wasb.csv").write_text("Frame,Visibility,X,Y\n")
650
651     lines = [
652         "some startup noise",
653         "1234/5000 (24%) | 2.3 win/s | elapsed 45s | ETA 120s",
654         "2500/5000 (50%)",
655         "final line",
656     ]
657
658     class _FakeStdout:
659         def __init__(self, data):
660             self._data = data
661             self._i = 0
662
663         def __iter__(self):
664             return self
665
666         def __next__(self):
667             if self._i >= len(self._data):
668                 raise StopIteration
669             val = self._data[self._i]
670             self._i += 1
671             return val + "\n"
672
673     received = []
674
675     def _fake_popen(argv, cwd=None, stdout=None, stderr=None, text=None):
676         # simulate the proc
677         proc = types.SimpleNamespace(
678             stdout=_FakeStdout(lines),
679             returncode=0,
680             poll=lambda: 0,
681             communicate=lambda: ("", ""),
682             kill=lambda: None,
683         )
684         return proc
685
686     with patch.object(r, "_copy_infer_script", return_value=True), patch(
687         "backend.pipeline.detectors.native_wasb_runner.subprocess.Popen",
688         side_effect=_fake_popen,
689     ), patch.object(r, "_rm_rf"), caplog.at_level(logging.INFO):
690         out = r.run_predict(
691             r"C:\v\clip.mp4",
692             str(tmp_path),
693             video_id="abc123abc123",
694             progress=received.append,
695         )
696
697     assert out is not None
698     # Only the two matching % lines should have produced callbacks (exact format from
parse)
699     assert received == ["detector 24% (1234/5000)", "detector 50% (2500/5000)"]
700     # The runner still logged the lines (via logger.info inside the loop)

```

```
701         assert any("1234/5000 (24%" in rec.getMessage() for rec in caplog.records)
702
```

49. assistant (2026-07-09T06:33:45.217Z)

This is the exact streaming test I need to extend. My runner change (merge stderr?stdout, ``proc.wait()`` instead of ``communicate()``) means I'll update this test's fake to accept ``**kwargs`` and expose ``wait()``. Before changing signatures, let me confirm nothing else in the codebase calls the two report helpers, and re-read the precise `cmd_infer` region I'll edit.

50. user (2026-07-09T06:33:47.911Z)

```
backend\worker\__main__.py:99:def _write_report_json(
backend\worker\__main__.py:125:def _serialize_and_push_outputs(
backend\worker\__main__.py:145:     _write_report_json(
backend\worker\__main__.py:318:         _serialize_and_push_outputs(
backend\worker\__main__.py:480:     pushed = _serialize_and_push_outputs(
tests\test_worker.py:73:def test_write_report_json(tmp_path):
tests\test_worker.py:75:     wmain._write_report_json(
```

51. assistant (2026-07-09T06:36:27.072Z)

I now have full understanding of all three gaps and the test patterns. Let me confirm my root-cause analysis and plan, then implement.

****Root causes:****

- ****GAP 3**** is two compounding bugs: (1) ``wasb_infer.py`` never configures logging, so its ``logger.info`` progress lines sit below `root=WARNING` and are suppressed; (2) even when emitted they go to stderr, but ``native_wasb_runner.py`` reads only stdout (and draining stdout while stderr's pipe fills risks a deadlock).
- ****GAP 1/2**** are in ``cmd_infer``: no per-stage logging on success, no timing, and neither timings nor validator verdicts reach the pushed ``report.json``.

Let me mark this phase and start implementing. First, the emission fix in ``wasb_infer.py``.

52. user (2026-07-09T06:36:28.209Z)

Chapter marked: "Implement observability fixes". Continue your current work.

53. user (2026-07-09T06:36:33.576Z)

The file `C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\detectors\wasb_infer.py` has been updated successfully. (file state is current in your context ? no need to Read it back)

54. user (2026-07-09T06:36:43.242Z)

The file `C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\detectors\wasb_infer.py` has been updated successfully. (file state is current in your context ? no need to Read it back)

55. assistant (2026-07-09T06:36:47.708Z)

Now the capture fix in ``native_wasb_runner.py``: add ``deque``, merge stderr?stdout, use ``proc.wait()``, and keep a bounded tail for the failure message.

56. user (2026-07-09T06:36:49.992Z)

The file `C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\detectors\native_wasb_runner.py` has been updated successfully. (file state is current in your context ? no need to Read it back)

57. user (2026-07-09T06:37:12.073Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\detectors\native_wasb_runner.py has been updated successfully. (file state is current in your context ? no need to Read it back)

58. assistant (2026-07-09T06:37:18.279Z)

Now the worker `\_\_main\_\_.py`. First, extend `\_write\_report\_json` to persist `timings` + `validation`:

59. user (2026-07-09T06:37:26.669Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

60. assistant (2026-07-09T06:37:34.216Z)

Now `\_serialize\_and\_push\_outputs` ? thread the two new args through:

61. user (2026-07-09T06:37:44.362Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

62. assistant (2026-07-09T06:37:52.823Z)

Now add the validation-summary helper just before `cmd\_infer`:

63. user (2026-07-09T06:38:00.852Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

64. assistant (2026-07-09T06:38:18.174Z)

Now the core `cmd\_infer` flow ? checksum timing + log, validation timing + PASS/FAIL log, shared `validation\_payload`, and threading both into `\_fail`:

65. user (2026-07-09T06:38:36.574Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

66. assistant (2026-07-09T06:38:45.251Z)

Now thread `timings` + `validation\_payload` into the `\_fail` closure's push:

67. user (2026-07-09T06:38:48.553Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

68. assistant (2026-07-09T06:39:04.409Z)

Now add segmentation timing + the per-stage log line:

69. user (2026-07-09T06:39:08.765Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-

indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

70. assistant (2026-07-09T06:39:18.115Z)

Finally, thread both into the success-path push:

71. user (2026-07-09T06:39:21.535Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

72. assistant (2026-07-09T06:39:30.985Z)

Let me re-read the modified `cmd\_infer` end-to-end to verify the closure ordering and that nothing is broken:

73. user (2026-07-09T06:39:31.724Z)

```
288         with no numeric detail and to the lightweight fakes used in tests (no ``details``
attribute)."""
289         name = getattr(r, "validator_name", "?")
290         ok = bool(getattr(r, "passed", False))
291         details = getattr(r, "details", None) or {}
292         hint = ""
293         for k in _VALIDATOR_SCORE_KEYS:
294             if k in details:
295                 hint = f"({k}={details[k]})"
296                 break
297         return f"{name}={ 'ok' if ok else 'FAIL' }{hint}"
298
299
300     def cmd_infer(args: argparse.Namespace) -> int:
301         config = load_config(args.config) if args.config else load_config()
302         _apply_overrides(config, args.set)
303         if not _run_startup_checks(config):
304             return 2
305
306         # M6: a cloud compute_target DISPATCHES to a Cloud Run Job instead of running in-
process.
307         # DEFAULT-OFF: get_compute_backend returns None for ""/local_gpu/cpu_mock ? fall
through to the
308         # unchanged inline path below. The dispatched container runs THIS cmd_infer inline
(the
309         # dispatcher forces compute_target=local_gpu), so it never recursively re-
dispatches.
310         compute = get_compute_backend(config)
311         if compute is not None:
312             return _dispatch_remote(compute, args)
313
314         storage_cfg = config.storage.model_dump()
315
316         scratch = args.scratch or tempfile.mkdtemp(prefix="rally-worker-")
317         os.makedirs(scratch, exist_ok=True)
318         config.output.output_dir = scratch
319
320         # 1. PULL ? local:// / bare path = identity passthrough; s3:// / gcs:// = download
to scratch.
321         local_video = storage.fetch(
322             args.video_uri, os.path.join(scratch, "in.mp4"), config=storage_cfg
323         )
324         print(f"[worker] pulled {args.video_uri} -> {local_video}")
325
```

```

326         # Per-stage wall-clock (#534 follow-up, 2026-07-09 live CUJ): each stage is timed
with
327         # perf_counter and persisted into report.json, so a runlog shows time-per-stage
instead of the
328         # earlier black hole (a ~10-min segmentation recorded nothing anywhere). Built
incrementally so
329         # a mid-run _fail() still reports the stages that DID run.
330         timings: dict[str, float] = {}
331
332         # 2. IDENTITY ? whole-file MD5 (the labels?video join key; pulled bytes must be
byte-identical).
333         _t = time.perf_counter()
334         video_id = compute_video_id(local_video)
335         timings["checksum_s"] = round(time.perf_counter() - _t, 3)
336         print(f"[worker] checksum: video_id={video_id} ({timings['checksum_s']:.1f}s)")
337
338         # 3. VALIDATE (same registry as the main CLI) ? UNLESS the dispatcher already did.
339         # On the cloud-serving path the CONTROL-PLANE is the authoritative validation gate:
it runs the
340         # SAME validators with ITS resolved config and only dispatches on PASS. Re-
validating here is
341         # redundant, decodes the whole video a SECOND time, and is outright WRONG when the
worker's config
342         # resolves different thresholds than the control-plane's ? observed 2026-07-07: a
5312x2988 clip
343         # (sharpness 11.9) PASSED the control-plane at validation.min_blur_threshold=8.0,
then the worker
344         # REJECTED it at its own default 20.0, so a SUCCESSFUL GPU dispatch failed after the
fact.
345         # --skip-validation (set by the dispatcher in orchestration._maybe_dispatch_remote)
skips it; the
346         # direct ``worker infer`` CLI has no control-plane, so it still validates by
default.
347         #
348         # When it DOES run, log the verdict loudly (PASS or FAIL): on the cloud-serving path
349         # phase_placement.validation=cloud_run_l4 means the worker really does run the full
suite, but a
350         # PASS previously logged nothing at all ? users reasonably concluded validation was
skipped.
351         if getattr(args, "skip_validation", False):
352             print(
353                 "[worker] validation SKIPPED (--skip-validation): the dispatcher is the
authoritative "
354                 "gate and already validated this input."
355             )
356             val_results, val_context = [], {}
357             timings["validation_s"] = 0.0
358         else:
359             n_checks = len(getattr(validator_registry, "_validators", []))
360             print(f"[worker] validating ({n_checks} checks?)")
361             _t = time.perf_counter()
362             val_results, val_context = validator_registry.run_all_with_context(
363                 local_video, config
364             )
365             timings["validation_s"] = round(time.perf_counter() - _t, 3)
366             verdict = "FAILED" if any(not r.passed for r in val_results) else "PASSED"
367             summary = " ".join(_val_token(r) for r in val_results) or "(no checks ran)"
368             print(
369                 f"[worker] validation {verdict} in {timings['validation_s']:.1f}s ?
{summary}"
370             )
371         # The per-validator verdicts (name/passed/message/details) ? stored BOTH in the
ephemeral worker
372         # DB (below) and, now, in the pushed report.json (so a runlog captures validator
outcomes).

```

```

373     validation_payload = [r.model_dump() for r in val_results]
374     failures = [r for r in val_results if not r.passed]
375     db = Database(os.path.join(scratch, config.output.db_name))
376     db.add_video(
377         video_id,
378         local_video,
379         "failed" if failures else "processed",
380         validation_payload,
381     )
382
383     def _fail(rc: int) -> int:
384         # A worker-side FAILURE (validation, unknown segmenter, or a detector
385         # segmenter surfaced as a Failure) must NOT masquerade as a 0-rally "success".
386         # a status="failed" report.json/intervals.csv (0 segments) so the control
387         # truth (submit-cache probe, restart-recovery) and the alpha-user result surface
388         # explicit FAILURE ? the wf294 silent-failure fix. Best-effort: a push hiccup
389         # mask the original failure. THEN the M6 done-marker (real rc + status) so a
390         # dispatcher's bucket-poll terminates FAST instead of hanging until its budget
391         # expires.
392         try:
393             _serialize_and_push_outputs(
394                 scratch,
395                 args.out_uri,
396                 video_id,
397                 [],
398                 "failed",
399                 storage_cfg,
400                 str(getattr(args, "video_uri", "") or ""),
401                 timings=timings,
402                 validation=validation_payload,
403             )
404         except Exception as e: # never mask the real failure on an artifact-push hiccup
405             logger.warning("[worker] could not push failure artifacts: %s", e)
406             # DEFAULT-OFF: no --done-uri ? no marker written (the unchanged local run).
407             if getattr(args, "done_uri", None):
408                 _write_done_marker(
409                     args.done_uri, video_id, rc, 0, "failed", storage_cfg, scratch
410                 )
411             return rc
412
413     segments: List[dict] = []
414     segmenter_actual = None
415     segmentation_ran = False
416     status = "failed"
417     try:
418         if failures:
419             print(
420                 "[worker] validation FAILED: "
421                 + "; ".join(f"{r.validator_name}: {r.message}" for r in failures)
422             )
423             return _fail(2)
424         seg_name = args.segmenter or config.indexing.default_segmenter
425         segmenter_cls = segmenter_registry.get(seg_name)
426         if not segmenter_cls:
427             print(
428                 f"[worker] unknown segmenter: {seg_name!r} "
429                 f"(have {list(segmenter_registry.list_available())})"
430             )
431             return _fail(2)

```

```

431     segmenter_actual = seg_name
432     out_local = os.path.join(scratch, "outputs", video_id) # for sidecar +
telemetry
433     os.makedirs(out_local, exist_ok=True)
434
435     # Live telemetry black-box for L4 detector runs (#534): velocity + GPU/CPU/RAM
samples.
436     # Written under outputs/ so the final push will mirror it; snapshots driven from
detector %.
437     telem: Optional[RunTelemetry] = None
438     try:
439         telem = RunTelemetry(
440             os.path.join(out_local, "run_telemetry.jsonl"),
441             label=f"l4-infer-{video_id[:8]}",
442         )
443     except Exception: # noqa: BLE001 - telemetry must never break a production run
444         telem = None
445
446     out_prefix = getattr(args, "out_uri", None)
447     if out_prefix:
448         out_prefix = out_prefix.rstrip("/")
449
450     # Configurable sampling for sidecar upload + telemetry snapshots (#534).
451     # Local sidecar file is always kept up-to-date (cheap). Only the costly remote
put
452     # and nvidia-smi-using snapshot are rate-limited. 0/negative = every callback.
453     sample_sec = 15.0
454     try:
455         idx_cfg = getattr(config, "indexing", None)
456         if idx_cfg is not None:
457             raw = getattr(idx_cfg, "detector_progress_sample_sec", 15.0)
458             sample_sec = float(raw) if raw is not None else 15.0
459     except Exception: # noqa: BLE001
460         sample_sec = 15.0
461
462     last_sample_t = [0.0]
463
464     def _progress(s: str) -> None:
465         print(f"[worker] progress: {s}")
466         # Write sidecar for cloud/L4 poll to surface live detector progress in job
status (#534).
467         # CP will read outputs/<video_id>/_detector_progress.json during poll.
468         # The *local* file is always written so the latest state is available in
scratch.
469         now = time.time()
470         do_sample = sample_sec <= 0 or (now - last_sample_t[0] >= sample_sec)
471
472         pfile = os.path.join(out_local, f"{video_id}_detector_progress.json")
473         try:
474             with open(pfile, "w") as pf:
475                 json.dump({"msg": s, "t": now}, pf)
476         except Exception:
477             pass # best effort
478
479         heavy_done = False
480         # LIVE push to storage (the part visible to CP poll during the run).
481         if out_prefix and do_sample:
482             try:
483                 storage.put(
484                     pfile,
485                     f"{out_prefix}/outputs/{video_id}/{os.path.basename(pfile)}",
486                     config=storage_cfg,
487                 )
488                 heavy_done = True
489             except Exception:

```



```

490             pass # best-effort live heartbeat
491
492     # Drive RunTelemetry snapshots (includes nvidia-smi + file rewrite) only on
sample.
493     if telem and s and "detector" in (s or "") and do_sample:
494         try:
495             m = re.search(r"(\d+)/(\d+)", s)
496             done = int(m.group(1)) if m else None
497             total = int(m.group(2)) if m else None
498             telem.snapshot(done=done, total=total, phase="detector")
499             heavy_done = True
500         except Exception: # noqa: BLE001
501             pass
502
503     if heavy_done:
504         last_sample_t[0] = now
505
506     seg_t = time.perf_counter()
507     result = segmenter_cls(db, config).process_video(
508         video_id,
509         local_video,
510         max_frames=args.max_frames,
511         progress=_progress,
512     )
513     segmentation_ran = True
514     timings["segmentation_s"] = round(time.perf_counter() - seg_t, 3)
515     if isinstance(result, Failure):
516         print(
517             f"[worker] segmentation ({segmenter_actual}) FAILED after "

```

74. assistant (2026-07-09T06:39:40.753Z)

Let me verify the remainder of the segmentation block and the tail:

75. user (2026-07-09T06:39:42.135Z)

```

515     if isinstance(result, Failure):
516         print(
517             f"[worker] segmentation ({segmenter_actual}) FAILED after "
518             f"{timings['segmentation_s']:.1f}s ? {result.message}"
519         )
520         return _fail(3)
521     segments = result.value if hasattr(result, "value") else result
522     status = "success"
523     print(
524         f"[worker] segmentation ({segmenter_actual}) done in "
525         f"{timings['segmentation_s']:.1f}s ? {len(segments)} rallies"
526     )
527     # Loud, never-silent diagnostic: a 0-segment success is almost always a
misconfig, not an
528     # empty match ? the cold-start failure mode (the substrate detector yielded no
usable
529     # trajectories). Explain it so a run is analysable from the log alone (re-run -v
for more).
530     if not segments:
531         detector_impl = (
532             os.environ.get("RALLY_DETECTOR_IMPL")
533             or getattr(config.indexing, "detector_impl", None)
534             or "auto"
535         )
536         logger.warning(
537             "infer produced 0 segments for video_id=%s (segmenter=%s,
detector_impl=%s). "
538             "The detector substrate yielded no usable trajectories/windows ? check
the detector "

```

```

539         "logs above; common causes: a missing ${AI_PROVIDER}_API_KEY with "
540         "skip_ai_handoff=False (the segmenter bails BEFORE detection ? set the
key or "
541         "indexing.skip_ai_handoff=true; cloud-serving now fails this at
startup), native "
542         "runner UNAVAILABLE (weights/repo/WASB_PYTHON env), the WASB env
produced no "
543         "detections, or the input resolution differs from the tuned proxy
resolution. "
544         "Re-run with -v for the native command + frame counts.",
545         video_id,
546         segmenter_actual,
547         detector_impl,
548     )
549     finally:
550         # Best-effort per-video observability report (never raises) ? parity with the
main CLI.
551         emit_indexer_report(
552             db=db,
553             video_id=video_id,
554             video_path=local_video,
555             config=config,
556             val_results=val_results,
557             val_context=val_context,
558             segmenter_requested=args.segmenter,
559             segmenter_actual=segmenter_actual,
560             was_reingest=False,
561             segmentation_ran=segmentation_ran,
562         )
563
564     # 5. SERIALIZE + 6. PUSH (outputs/<video_id>/?). idempotent on video_id.
565     pushed = _serialize_and_push_outputs(
566         scratch,
567         args.out_uri,
568         video_id,
569         segments,
570         status,
571         storage_cfg,
572         str(getattr(args, "video_uri", "") or ""),
573         timings=timings,
574         validation=validation_payload,

```

76. assistant (2026-07-09T06:41:41.558Z)

The worker ``cmd_infer`` flow is correct ? ``timings`` and ``validation_payload`` are defined before the ``_fail`` closure and threaded into every push. Now let me check how ``wasb_infer.py`` is imported by existing tests and whether it has heavy module-level imports (to pick the right test approach for the emission fix).

77. user (2026-07-09T06:41:43.814Z)

```

49:import argparse
50:import csv
51:import glob
52:import json
53:import logging
54:import os
55:import os.path as osp
56:import sys
57:from collections import defaultdict
58:from contextlib import contextmanager

```

78. user (2026-07-09T06:41:44.981Z)

```
28:from backend.pipeline.detectors import wasb_infer as wi
339:         spec = importlib.util.spec_from_file_location("wasb_infer_staged", staged)
```

79. assistant (2026-07-09T06:42:57.360Z)

Module-level imports are stdlib-only, so `wasb\_infer` imports cleanly. To make the emission fix independently testable (and to avoid mutating global logging inside `main()` untestably), let me extract the logging config into a small `\_configure\_logging()` helper ? mirroring the worker's tested `\_setup\_logging` pattern.

80. user (2026-07-09T06:43:06.244Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\detectors\wasb_infer.py has been updated successfully. (file state is current in your context ? no need to Read it back)

81. assistant (2026-07-09T06:43:17.924Z)

Now the tests. First, update the existing streaming test's fake to match the new Popen call (`bufsize`, `wait()`):

82. user (2026-07-09T06:43:21.858Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_native_wasb_runner.py has been updated successfully. (file state is current in your context ? no need to Read it back)

83. assistant (2026-07-09T06:43:34.600Z)

Now append the new native-runner tests (stderr-merge, failure-tail) and the `wasb\_infer` emission test:

84. user (2026-07-09T06:43:54.400Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_native_wasb_runner.py has been updated successfully. (file state is current in your context ? no need to Read it back)

85. assistant (2026-07-09T06:44:05.255Z)

Now the worker tests. First add the `types` import:

86. user (2026-07-09T06:44:07.265Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

87. assistant (2026-07-09T06:44:10.099Z)

Now add the pure-helper tests (report serialization + `\_val\_token`) after the existing `test\_write\_report\_json`:

88. user (2026-07-09T06:44:23.470Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

89. assistant (2026-07-09T06:45:04.435Z)

Now the two end-to-end `cmd\_infer` tests (PASS logs + report shape; FAIL timings/logging), placed after `test\_cmd\_infer\_local\_roundtrip`:

90. user (2026-07-09T06:45:20.781Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

91. assistant (2026-07-09T06:45:37.235Z)

All edits are in. Now let me run the verification gates. First, a quick environment probe (Python, ruff, mypy, pytest versions) so I know what's installed.

92. user (2026-07-09T06:45:40.826Z)

```
Python 3.13.13
--- ruff ---
ruff 0.15.16
--- mypy ---
mypy 2.1.0 (compiled: yes)
--- pytest ---
pytest 8.4.2
```

93. assistant (2026-07-09T06:45:50.249Z)

Toolchain matches CI exactly (ruff 0.15.16, mypy 2.1.0) ? no installs needed. Let me run the directly-affected tests first for fast feedback.

94. user (2026-07-09T06:45:59.740Z)

```
..... [ 70%]
..... [100%]
102 passed in 6.20s
```

95. assistant (2026-07-09T06:46:08.988Z)

All 102 pass. Now ruff and mypy:

96. user (2026-07-09T06:46:09.921Z)

All checks passed!

97. user (2026-07-09T06:46:19.068Z)

```
backend\annotations\base.py:88: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\sports\__init__.py:9: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\pipeline\validators\base.py:123: note: By default the bodies of untyped functions are
not checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 150 source files
```

98. assistant (2026-07-09T06:46:39.042Z)

Ruff clean, mypy clean (150 files, only informational notes). Now the full suite with coverage. Let me first confirm `run\_all\_tests.py` passes through the `--cov` flags so I invoke it correctly.

99. user (2026-07-09T06:46:40.170Z)

```
1      #!/usr/bin/env python
2      """Run the FULL test suite for the badminton-highlight-indexer.
3
4      A single, discoverable entry point so you (or a reviewer / CI) can verify the
5      whole suite before opening a PR:
6
```

```

7         python run_all_tests.py                # full suite, quiet
8         python run_all_tests.py -v             # verbose (passthrough to pytest)
9         python run_all_tests.py -k config      # filter by keyword (passthrough)
10        python run_all_tests.py --cov=backend --cov-report=term-missing # coverage
11
12    Anything after the script name is passed straight through to pytest.
13
14    Notes
15    -----
16    - The suite is fully mocked (no GPU, no network, no WSL) and runs offline.
17    - The observability **reporting** tests skip cleanly if the optional ``obsreport``
18      package isn't installed (they use ``pytest.importorskip``). To exercise them too:
19      pip install -e ../sports-obsreport --no-build-isolation
20    - Exit code is pytest's own, so this is CI / pre-PR friendly (0 = all green).
21    ""
22
23    from __future__ import annotations
24
25    import importlib.util
26    import subprocess
27    import sys
28
29
30    def _installed(module: str) -> bool:
31        try:
32            return importlib.util.find_spec(module) is not None
33        except (ImportError, ValueError):
34            return False
35
36
37    def main() -> int:
38        args = list(sys.argv[1:])
39        # Default to quiet unless the caller chose a verbosity.
40        if not any(a in ("-q", "-v", "-vv", "-vvv", "--quiet", "--verbose") for a in args):
41            args = ["-q", *args]
42        cmd = [sys.executable, "-m", "pytest", *args]
43
44        obs = (
45            "available"
46            if _installed("obsreport")
47            else "NOT installed ? reporting tests will SKIP"
48        )
49        bar = "=" * 72
50        print(bar)
51        print(" Full test suite: " + " ".join(cmd))
52        print(f" obsreport (optional reporting dep): {obs}")
53        print(bar, flush=True)
54
55        rc = subprocess.run(cmd).returncode
56
57        print(bar)
58        print(
59            " RESULT: "
60            + ("PASS - all tests green" if rc == 0 else f"FAIL - pytest exit {rc}")
61        )
62        print(bar)
63        return rc
64
65
66    if __name__ == "__main__":
67        raise SystemExit(main())
68

```

`run\_all\_tests.py` passes args straight through to pytest. Running the full suite with coverage now (this is the pre-PR gate):

101. user (2026-07-09T06:53:51.216Z)

backend\pipeline\validators\format.py	36	16	56%
backend\pipeline\validators\frame_sampler.py	54	1	98%
backend\pipeline\validators\pts_continuity.py	56	39	30%
backend\pipeline\validators\relevance.py	51	23	55%
backend\pipeline\validators\resolution_fps.py	50	23	54%
backend\pipeline\validators\scene_cut.py	57	38	33%
backend\providers__init__.py	16	1	94%
backend\providers\base.py	9	1	89%
backend\providers\gemini.py	137	21	85%
backend\reporting__init__.py	26	2	92%
backend\reporting\harvest.py	111	8	93%
backend\reporting\run_signals.py	14	0	100%
backend\results.py	27	0	100%
backend\server_runtime.py	57	2	96%
backend\sports__init__.py	17	1	94%
backend\sports\badminton.py	13	1	92%
backend\sports\base.py	12	3	75%
backend\storage__init__.py	128	16	88%
backend\storage\base.py	16	1	94%
backend\storage\capabilities.py	66	11	83%
backend\storage\gcs.py	80	3	96%
backend\storage\gdrive.py	61	3	95%
backend\storage\gdrive_api.py	105	6	94%
backend\storage\local.py	49	3	94%
backend\storage\ref.py	49	0	100%
backend\storage\s3.py	47	5	89%
backend\tools__init__.py	0	0	100%
backend\tools\export.py	75	2	97%
backend\tools\rally_slicer.py	105	47	55%
backend\utils__init__.py	0	0	100%
backend\utils\annotation_proxy.py	64	11	83%
backend\utils\app_home.py	40	0	100%
backend\utils\ffmpeg.py	47	3	94%
backend\utils\freospace.py	32	2	94%
backend\utils\geometry.py	29	0	100%
backend\utils\hashing.py	9	0	100%
backend\utils\logging_setup.py	11	0	100%
backend\utils\paths.py	44	3	93%
backend\utils\redact.py	16	0	100%
backend\utils\run_telemetry.py	126	12	90%
backend\utils\single_instance.py	43	11	74%
backend\utils\video.py	34	3	91%
backend\utils\workspace.py	49	0	100%
backend\worker__init__.py	0	0	100%
backend\worker__main__.py	457	43	91%
training__init__.py	0	0	100%
training\gen0__init__.py	0	0	100%
training\gen0\harness.py	130	9	93%

TOTAL 13582 2114 84%

Required test coverage of 70% reached. Total coverage: 84.44%

1965 passed, 1 skipped, 57 warnings in 417.23s (0:06:57)

=====
RESULT: PASS - all tests green
=====

102. assistant (2026-07-09T06:54:28.020Z)

Full suite: **1965 passed, 1 skipped**, coverage **84.44%** (?70% floor), worker `\_\_main\_\_.py` at 91%. No `test\_storage.py` failures ? #548 (already on origin/master) resolved them. ruff and

mypy clean.

One backward-compat check before I commit: I added `timings`/`validation` keys to `report.json`. Let me verify the control-plane's reader doesn't reject unknown keys (e.g. a Pydantic model with `extra="forbid"`).

103. user (2026-07-09T06:54:30.399Z)

```
111:     if report.get("output_dir"):
314:         uri, os.path.join(td, ref.name or "report.json"), config=storage_config
360: def probe_process_cache(db, cfg, video_id: str) -> bool:
381:     report_uri = f"{output_root}/outputs/{video_id}/report.json"
385:     return bool(report and report.get("status") == "success")
395: def cached_process_result(
399:     the DB first, then the bucket (`outputs/<md5>/report.json` with `status=success`
429:     report_uri = f"{outputs_uri}/report.json"
433:     if not report or report.get("status") != "success":
441:         video_path = str(report.get("video_uri") or "") or _first_object_under(
493:     if not probe_process_cache(db, cfg, vid):
495:     return cached_process_result(db, cfg, vid, "") is not None
1685:     result.report.get("message")
1691:     result.report.get("download_url")
1704:     "filepath": str(result.report.get("reel_uri") or result.outputs_uri or ""),
1707:     "reel_uri": str(result.report.get("reel_uri") or ""),
1708:     "filename": str(result.report.get("filename") or out_name),
1709:     "rendition": str(result.report.get("rendition") or rendition),
```

104. assistant (2026-07-09T06:54:43.594Z)

Confirmed? the CP reads `report.json` via plain `report.get(...)` , so the additive `timings`/`validation` keys are ignored by existing readers. Backward compatible.

Let me self-review the full diff before committing.

105. user (2026-07-09T06:54:49.901Z)

```
backend/pipeline/detectors/native_wasb_runner.py | 57 +++++--
backend/pipeline/detectors/wasb_infer.py         | 22 +++
backend/worker/__main__.py                       | 135 ++++++-----
tests/test_native_wasb_runner.py                 | 110 ++++++-----
tests/test_worker.py                             | 172 ++++++-----
5 files changed, 450 insertions(+), 46 deletions(-)
===== FULL DIFF: detectors =====
diff --git a/backend/pipeline/detectors/native_wasb_runner.py
b/backend/pipeline/detectors/native_wasb_runner.py
index 81ebbfd..1d80828 100644
--- a/backend/pipeline/detectors/native_wasb_runner.py
+++ b/backend/pipeline/detectors/native_wasb_runner.py
@@ -39,6 +39,7 @@ import os
 import re
 import shutil
 import subprocess
+from collections import deque
 from dataclasses import dataclass
 from typing import Callable, List, Optional, Tuple

@@ -369,41 +370,57 @@ class NativeWasbRunner(DetectorRunner):
     # Fast path: capture at end (existing behaviour, byte-identical for tests).
     res = self._run(argv, cwd=self._repo_src())
 else:
     # Live progress path (#534): stream stdout so we can parse the periodic
     # "done/total (pct%) ..." lines already emitted by wasb_infer.py and call
     # the callback. This gives the SPA a real heartbeat during the long GPU pass
     # instead of a spinner. Still respects timeout.
```

```

+         # Live progress path (#534): stream the child's output so we can parse the
+         # periodic "done/total (pct%) ..." lines wasb_infer emits and fire the callback
+         ?
+         # a real heartbeat during the long GPU pass instead of a spinner.
+         #
+         # MERGE stderr INTO stdout (stderr=STDOUT). wasb_infer routes its logging
(progress
+         # included) to a stream, and third-party libs may also write to stderr. Reading
+         # stdout alone would (a) miss any progress that lands on stderr and (b) risk a
+         # DEADLOCK: draining only stdout while the child fills its ~64 KB stderr pipe
blocks
+         # the child on the stderr write, which stops its stdout too ? the run then
hangs
+         # until timeout. One merged pipe removes both failure modes (the 2026-07-09
live CUJ
+         # produced zero heartbeat lines). bufsize=1 = line-buffered so the heartbeat is
+         # timely, not block-delayed. A bounded tail keeps the last lines for the
failure
+         # message, since there is no separate stderr to read on error.
proc = subprocess.Popen(
    argv,
    cwd=self._repo_src(),
    stdout=subprocess.PIPE,
-     stderr=subprocess.PIPE,
+     stderr=subprocess.STDOUT,
    text=True,
+     bufsize=1,
)
assert proc.stdout is not None
+ tail: "deque[str]" = deque(maxlen=200)
try:
    for line in proc.stdout:
        line = line.rstrip()
-         if line:
-             # Parse existing wasb_infer progress lines for detector heartbeat.
-             # Example: "1234/5000 (24%) | 2.3 win/s | elapsed 45s | ETA 120s"
-             m = re.search(r"(\d+)/(\d+)\s+\((\d+)\%\)", line)
-             if m:
-                 done, total, pct = m.groups()
-                 try:
-                     progress(f"detector {pct}% ({done}/{total})")
-                 except Exception: # noqa: BLE001
-                     pass # progress must never kill the run
-                 # Also log for the container/runlog.
-                 logger.info(line)
-             # Wait for completion (communicate for stderr/returncode).
-             _, stderr = proc.communicate()
+         if not line:
+             continue
+         tail.append(line)
+         # Parse wasb_infer progress lines for the detector heartbeat.
+         # Example: "1234/5000 (24%) | 2.3 win/s | elapsed 45s | ETA 120s"
+         m = re.search(r"(\d+)/(\d+)\s+\((\d+)\%\)", line)
+         if m:
+             done, total, pct = m.groups()
+             try:
+                 progress(f"detector {pct}% ({done}/{total})")
+             except Exception: # noqa: BLE001
+                 pass # progress must never kill the run
+             # Mirror to the container/runlog.
+             logger.info(line)
+         proc.wait()
res = subprocess.CompletedProcess(
    args=argv, returncode=proc.returncode, stdout="", stderr=stderr or ""
+     args=argv,

```



```

+                 returncode=proc.returncode,
+                 stdout="",
+                 stderr="\n".join(tail),
+             )
+         finally:
+             if proc.poll() is None:
+                 proc.kill()
+                 proc.wait()
+     except FileNotFoundError:
+         logger.error(
+             "python binary not found: %r (set WASB_PYTHON).", self.cfg.python_bin
diff --git a/backend/pipeline/detectors/wasb_infer.py b/backend/pipeline/detectors/wasb_infer.py
index ef43816..13c38a4 100644
--- a/backend/pipeline/detectors/wasb_infer.py
+++ b/backend/pipeline/detectors/wasb_infer.py
@@ -53,6 +53,7 @@ import json
import logging
import os
import os.path as osp
+import sys
from collections import defaultdict
from contextlib import contextmanager

@@ -1247,7 +1248,28 @@ def run_video_nvdec_resident(
    return _tracker_pass_and_write(cfg, det_by_fid, out_csv)

+def _configure_logging() -> None:
+    """Own logging config for this SUBPROCESS entrypoint (run as ``python wasb_infer.py ?`` by
+    the
+    + WSL / native runners). Without it the module logger inherits the root default (WARNING) and
+    + EVERY ``logger.info()`` below is silently dropped ? including the periodic "done/total
+    (pct%)"
+    + progress lines the native runner parses for the #534 live detector heartbeat. That is
+    exactly
+    + why a ~10-min L4 pass produced zero progress lines, zero telemetry snapshots, and no
+    + detector_progress sidecar (2026-07-09 live CUJ).
+    +
+    + Route to STDOUT ? the stream ``NativeWasbRunner`` streams + parses ? and ``force=True`` so
+    a
+    + library that pre-configured logging can't leave us pinned at WARNING. ``StreamHandler``
+    flushes
+    + per record, so the heartbeat is timely rather than block-buffered."""
+    logging.basicConfig(
+        level=logging.INFO,
+        format="%(asctime)s [%(levelname)s] %(name)s: %(message)s",
+        datefmt="%H:%M:%S",
+        stream=sys.stdout,
+        force=True,
+    )
+
+def main():
+    _configure_logging()
+    ap = argparse.ArgumentParser()
+    ap.add_argument("--frames_dir", help="folder of pre-extracted frames")
+    ap.add_argument(

```

106. assistant (2026-07-09T06:54:56.474Z)

The detector diffs are clean. Now the worker diff:

107. user (2026-07-09T06:54:57.494Z)

```
diff --git a/backend/worker/__main__.py b/backend/worker/__main__.py
```

```

index faccb9f..1692d84 100644
--- a/backend/worker/__main__.py
+++ b/backend/worker/__main__.py
@@ -97,29 +97,40 @@ def _write_intervals_csv(segments: List[dict], path: str) -> None:

def _write_report_json(
-    video_id: str, segments: List[dict], status: str, path: str, video_uri: str = ""
+    video_id: str,
+    segments: List[dict],
+    status: str,
+    path: str,
+    video_uri: str = "",
+    timings: Optional[dict] = None,
+    validation: Optional[List[dict]] = None,
) -> None:
-    with open(path, "w") as f:
-        json.dump(
+    payload: dict[str, Any] = {
+        "video_id": video_id,
+        "status": status,
+        "segment_count": len(segments),
+        # The SOURCE the detection ran on (CP-rebuild P3, #485): a restart-wiped
+        # control plane re-hydrating from outputs/<md5>/ restores the video row
+        # with its truthful path ? without this, only the segments recover.
+        "video_uri": video_uri,
+        "segments": [
+            {
-                "video_id": video_id,
-                "status": status,
-                "segment_count": len(segments),
-                # The SOURCE the detection ran on (CP-rebuild P3, #485): a restart-wiped
-                # control plane re-hydrating from outputs/<md5>/ restores the video row
-                # with its truthful path ? without this, only the segments recover.
-                "video_uri": video_uri,
-                "segments": [
-                    {
-                        "start": float(s.get("start_time", 0.0)),
-                        "end": float(s.get("end_time", 0.0)),
-                    }
-                    for s in segments
-                ],
-            },
-            f,
-            indent=2,
-        )
+                "start": float(s.get("start_time", 0.0)),
+                "end": float(s.get("end_time", 0.0)),
+            }
+            for s in segments
+        ],
+    },
+    f,
+    indent=2,
+    )
+        "start": float(s.get("start_time", 0.0)),
+        "end": float(s.get("end_time", 0.0)),
+    }
+    for s in segments
+    ],
+    }
+    # Observability (#534 follow-up, 2026-07-09 live CUJ): persist per-stage wall-clock AND the
+    # validator verdicts the worker actually computed. Previously both were invisible on a PASS
+    ?
+    # the results lived only in the ephemeral worker DB and never reached this pushed report,
+    so a
+    # passing validation looked skipped and a ~10-min segmentation recorded no timing anywhere.
+    if timings:
+        payload["timings"] = timings
+    if validation is not None:
+        payload["validation"] = validation
+    with open(path, "w") as f:
+        json.dump(payload, f, indent=2)

```

```
def _serialize_and_push_outputs(
@@ -130,6 +141,8 @@ def _serialize_and_push_outputs(
    status: str,
    storage_cfg: dict,
    video_uri: str = "",
+   timings: Optional[dict] = None,
+   validation: Optional[List[dict]] = None,
) -> List[str]:
    """Write intervals.csv + report.json under ``<scratch>/outputs/<video_id>/`` and push both
to
    ``<out_uri>/outputs/<video_id>/``; return the pushed URIs.
@@ -138,7 +151,10 @@ def _serialize_and_push_outputs(
    ``status="failed"`` report.json in the bucket ? the control plane's source of truth
    (`orchestration.cached_process_result` / `probe_process_cache` treat any
non-``success``
    report as "no usable result" ? retry, and the restart-recovery poll waits for report.json
to
-   EXIST) ? instead of a missing artifact or a false 0-segment ``"success"``.
+   EXIST) ? instead of a missing artifact or a false 0-segment ``"success"``.
+
+   ``timings`` (per-stage wall-clock) and ``validation`` (the per-validator verdicts) are
folded
+   into report.json for observability (see :func:`_write_report_json`)."""
    out_local = os.path.join(scratch, "outputs", video_id)
    os.makedirs(out_local, exist_ok=True)
    _write_intervals_csv(segments, os.path.join(out_local, "intervals.csv"))
@@ -148,6 +164,8 @@ def _serialize_and_push_outputs(
    status,
    os.path.join(out_local, "report.json"),
    video_uri=video_uri,
+   timings=timings,
+   validation=validation,
)
    out_prefix = out_uri.rstrip("/")
    pushed: List[str] = []
@@ -249,6 +267,36 @@ def _write_done_marker(
    logger.warning("[worker] could not write done-marker %s: %s", done_uri, e)

+# Headline metric key per validator (in ``ValidationResult.details``) for the one-line
validation
+# summary log. The FULL details dict is persisted in report.json; this is only the at-a-glance
+# score for the human-readable ``[worker] validation ?`` line. First key present wins.
+_VALIDATOR_SCORE_KEYS = (
+   "avg_laplacian_variance", # blur_check ? sharpness
+   "avg_green_court_ratio", # relevance ? court coverage
+   "cuts_per_minute", # scene_cut
+   "max_keyframe_gap_seconds", # pts_continuity
+   "fps", # resolution_fps
+   "duration", # format_check
+)
+
+def _val_token(r: Any) -> str:
+   """One ``name=ok`` / ``name=FAIL(score)`` token for the validation-summary log line.
+
+   ``score`` is a best-effort headline metric pulled from ``details`` (see
+:data:`_VALIDATOR_SCORE_KEYS`) and omitted when absent ? so this is robust to a validator
+   with no numeric detail and to the lightweight fakes used in tests (no ``details``
attribute)."""
+   name = getattr(r, "validator_name", "?")
+   ok = bool(getattr(r, "passed", False))
+   details = getattr(r, "details", None) or {}
+   hint = ""
+   for k in VALIDATOR SCORE KEYS:
```

```

+         if k in details:
+             hint = f"({k}={details[k]})"
+             break
+         return f"{name}={'ok' if ok else 'FAIL'}{hint}"
+
+
+ def cmd_infer(args: argparse.Namespace) -> int:
+     config = load_config(args.config) if args.config else load_config()
+     _apply_overrides(config, args.set)
@@ -275,8 +323,17 @@ def cmd_infer(args: argparse.Namespace) -> int:
+ )
+     print(f"[worker] pulled {args.video_uri} -> {local_video}")
+
+ # Per-stage wall-clock (#534 follow-up, 2026-07-09 live CUJ): each stage is timed with
+ # perf_counter and persisted into report.json, so a runlog shows time-per-stage instead of
the
+ # earlier black hole (a ~10-min segmentation recorded nothing anywhere). Built
incrementally so
+ # a mid-run _fail() still reports the stages that DID run.
+     timings: dict[str, float] = {}
+
+     # 2. IDENTITY ? whole-file MD5 (the labels?video join key; pulled bytes must be byte-
identical).
+     _t = time.perf_counter()
+     video_id = compute_video_id(local_video)
+     timings["checksum_s"] = round(time.perf_counter() - _t, 3)
+     print(f"[worker] checksum: video_id={video_id} ({timings['checksum_s']:.1f}s)")
+
+     # 3. VALIDATE (same registry as the main CLI) ? UNLESS the dispatcher already did.
+     # On the cloud-serving path the CONTROL-PLANE is the authoritative validation gate: it runs
the
@@ -287,23 +344,40 @@ def cmd_infer(args: argparse.Namespace) -> int:
+     # REJECTED it at its own default 20.0, so a SUCCESSFUL GPU dispatch failed after the fact.
+     # --skip-validation (set by the dispatcher in orchestration._maybe_dispatch_remote) skips
it; the
+     # direct ``worker infer`` CLI has no control-plane, so it still validates by default.
+     #
+     # When it DOES run, log the verdict loudly (PASS or FAIL): on the cloud-serving path
+     # phase_placement.validation=cloud_run_l4 means the worker really does run the full suite,
but a
+     # PASS previously logged nothing at all ? users reasonably concluded validation was
skipped.
+     if getattr(args, "skip_validation", False):
+         print(
+             "[worker] validation SKIPPED (--skip-validation): the dispatcher is the
authoritative "
+             "gate and already validated this input."
+         )
+         val_results, val_context = [], {}
+     timings["validation_s"] = 0.0
+     else:
+         n_checks = len(getattr validator_registry, "_validators", []))
+         print(f"[worker] validating ({n_checks} checks)?")
+         _t = time.perf_counter()
+         val_results, val_context = validator_registry.run_all_with_context(
+             local_video, config
+         )
+         timings["validation_s"] = round(time.perf_counter() - _t, 3)
+         verdict = "FAILED" if any(not r.passed for r in val_results) else "PASSED"
+         summary = " ".join(_val_token(r) for r in val_results) or "(no checks ran)"
+         print(
+             f"[worker] validation {verdict} in {timings['validation_s']:.1f}s ? {summary}"
+         )
+     # The per-validator verdicts (name/passed/message/details) ? stored BOTH in the ephemeral
worker

```

```

+ # DB (below) and, now, in the pushed report.json (so a runlog captures validator outcomes).
+ validation_payload = [r.model_dump() for r in val_results]
failures = [r for r in val_results if not r.passed]
db = Database(os.path.join(scratch, config.output.db_name))
db.add_video(
    video_id,
    local_video,
    "failed" if failures else "processed",
-    [r.model_dump() for r in val_results],
+    validation_payload,
)

def _fail(rc: int) -> int:
@@ -323,6 +397,8 @@ def cmd_infer(args: argparse.Namespace) -> int:
    "failed",
    storage_cfg,
    str(getattr(args, "video_uri", "") or ""),
+    timings=timings,
+    validation=validation_payload,
)
except Exception as e: # never mask the real failure on an artifact-push hiccup
    logger.warning("[worker] could not push failure artifacts: %s", e)
@@ -427,6 +503,7 @@ def cmd_infer(args: argparse.Namespace) -> int:
    if heavy_done:
        last_sample_t[0] = now

+    seg_t = time.perf_counter()
    result = segmenter_cls(db, config).process_video(
        video_id,
        local_video,
@@ -434,11 +511,19 @@ def cmd_infer(args: argparse.Namespace) -> int:
        progress=_progress,
    )
    segmentation_ran = True
+    timings["segmentation_s"] = round(time.perf_counter() - seg_t, 3)
    if isinstance(result, Failure):
-        print(f"[worker] segmenter FAILED: {result.message}")
+        print(
+            f"[worker] segmentation ({segmenter_actual}) FAILED after "
+            f"{timings['segmentation_s']:.1f}s ? {result.message}"
+        )
        return _fail(3)
    segments = result.value if hasattr(result, "value") else result
    status = "success"
+    print(
+        f"[worker] segmentation ({segmenter_actual}) done in "
+        f"{timings['segmentation_s']:.1f}s ? {len(segments)} rallies"
+    )
    # Loud, never-silent diagnostic: a 0-segment success is almost always a misconfig, not
an
    # empty match ? the cold-start failure mode (the substrate detector yielded no usable
    # trajectories). Explain it so a run is analysable from the log alone (re-run -v for
more).
@@ -485,6 +570,8 @@ def cmd_infer(args: argparse.Namespace) -> int:
    status,
    storage_cfg,
    str(getattr(args, "video_uri", "") or ""),
+    timings=timings,
+    validation=validation_payload,
)

print(

```

The worker diff is clean and all three gaps are addressed. Everything verifies: 1965 passed, ruff + mypy clean, coverage 84.44%, backward-compatible report.json. Let me re-verify branch safety (concurrent-session protocol), then stage explicit paths and commit.

109. user (2026-07-09T06:55:30.559Z)

```
=== current branch (must be mine) ===
fix/worker-l4-infer-observability
=== origin/master now ===
00e8f92
=== commits on my branch vs origin/master (must be EMPTY pre-commit) ===
(end)
```

110. assistant (2026-07-09T06:55:38.395Z)

Branch safety confirmed: on my branch, origin/master unchanged at `00e8f92`, clean base. Staging explicit paths only (never `-A`) and committing:

111. user (2026-07-09T06:55:40.690Z)

```
=== staged ===
M backend/pipeline/detectors/native_wasb_runner.py
M backend/pipeline/detectors/wasb_infer.py
M backend/worker/__main__.py
M tests/test_native_wasb_runner.py
M tests/test_worker.py
?? "SESSION_HANDOFF (1).md"
=== confirm HEAD is mine ===
fix/worker-l4-infer-observability
```

112. assistant (2026-07-09T06:55:46.688Z)

Exactly my 5 files staged (the stray untracked handoff file is not). Committing:

113. user (2026-07-09T06:56:01.038Z)

```
[fix/worker-l4-infer-observability b3dff78] fix(worker): make the cloud-serving L4 infer path
observable
5 files changed, 450 insertions(+), 46 deletions(-)
```

114. assistant (2026-07-09T06:56:09.330Z)

Committed as `b3dff78` (5 files, +450/?46). Verifying only my commit rides on origin/master, then pushing:

115. user (2026-07-09T06:56:13.499Z)

```
=== commits vs origin/master (must be ONLY mine) ===
b3dff78 fix(worker): make the cloud-serving L4 infer path observable
=== push ===
remote:
remote: Create a pull request for 'fix/worker-l4-infer-observability' on GitHub by visiting:
remote:      https://github.com/Khelsutra/badminton-highlight-
indexer/pull/new/fix/worker-l4-infer-observability
remote:
branch 'fix/worker-l4-infer-observability' set up to track 'origin/fix/worker-l4-infer-
observability'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
* [new branch]      fix/worker-l4-infer-observability -> fix/worker-l4-infer-observability
```

116. assistant (2026-07-09T06:56:23.201Z)

Pushed cleanly with only my commit. Now opening the PR to master:

117. user (2026-07-09T06:56:44.981Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/550>

118. assistant (2026-07-09T06:57:03.723Z)

PR #550 is open: <https://github.com/Khelsutra/badminton-highlight-indexer/pull/550>

Let me record the one non-derivable fact for future sessions ? that the live-L4 heartbeat still needs an owner-driven CUJ to confirm end-to-end ? as a memory with an index pointer.

119. assistant (2026-07-09T06:57:19.282Z)

name: pr550-worker-l4-observability

description: "PR #550 (open) makes the cloud-serving L4 worker infer path observable; live-L4 heartbeat still needs an owner CUJ"

metadata:

type: project

PR #550 ? cloud-serving L4 infer observability (OPEN, 2026-07-09)

Branch `fix/worker-l4-infer-observability` cut from `origin/master @ 00e8f92`; commit `b3dff78`. Fixes three linked observability gaps on `backend/worker/\_\_\_main\_\_.py`'s infer path, surfaced by the

2026-07-09 live CUJ (video GX010155, deployed image `sha256:1838c07f?`, worker `rally-worker-dpssz`).

Because `phase\_placement.validation.segment.compile` are all `cloud\_run\_l4`, the worker runs the full

validator suite + segmentation but none of it was observable.

- **GAP 1 (validation invisible on PASS):** now logs `[worker] validating (N checks)?` + `[worker] validation PASSED/FAILED in Ys ? <name>=ok/FAIL(score)?`, and the per-validator verdicts are folded into the pushed `report.json` under `validation` (previously only in the ephemeral DB).

- **GAP 2 (no per-stage timing):** checksum/validation/segmentation each timed with `time.perf\_counter()`, logged, and persisted as a `timings` block in `report.json` (built incrementally so a mid-run `\_fail()` still records the stages that ran).

- **GAP 3 (#534 heartbeat never fired on native-L4):** TWO compounding bugs ?

(a) `wasb\_infer.py` never configured logging so its `logger.info` `done/total (pct%)` progress sat

below root=WARNING and was dropped; (b) even when emitted it goes to **stderr** but `NativeWasbRunner` read only **stdout** (and draining stdout while stderr's ~64 KB pipe fills risks

a deadlock). Fix: `wasb\_infer.\_configure\_logging()` ? INFO to **stdout**; runner merges **stderr?stdout** (line-buffered) + bounded failure tail. Callback now fires ?

`run\_telemetry.jsonl`

snapshots + `[worker] progress: heartbeat + <video\_id>\_detector\_progress.json` on GCS.

`report.json` changes are additive; CP reads it via `report.get(...)` (`probe\_process\_cache` / `cached\_process\_result`) so existing readers are unaffected.

Verified locally (Windows): `run\_all\_tests.py --cov=backend --cov=training --cov-fail-under=70`

? 1965 passed / 1 skipped, coverage 84.44% (worker `\_\_\_main\_\_.py` 91%); ruff 0.15.16 + mypy 2.1.0 clean.

No `test\_storage.py` failures ? #548 (already on master) fixed those on Windows.

OPEN follow-up (not verifiable offline): the live-L4 heartbeat end-to-end (real GPU + `wasb` env +

Cloud Run L4) needs an **owner-driven CUJ** to confirm the telemetry snapshots / progress sidecar /

stdout heartbeat actually land. The offline tests lock the emission config, the stdout parse+merge contract, and the report/timing serialization ? but not the real-hardware behavior. Related: [[ci-windows-coverage-gap]] (why the local windows run matters), [[live-deployment-state]] (deploy the image after merge to make this live).

120. user (2026-07-09T06:57:19.327Z)

File created successfully at: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\pr550-worker-l4-observability.md (file state is current in your context ? no need to Read it back)

121. user (2026-07-09T06:57:24.977Z)

```
1      # Memory index
2
3      - ? **Start here (GPU-resident rollout):** [session-handoff-gpu-resident.md](session-
handoff-gpu-resident.md) ? #506 implemented, PR #507 open; next: on-box parity gate + 15-clip
golden gate.
4      - ? **LIVE deploy (2026-07-09):** [live-deployment-state.md](live-deployment-state.md) ?
CP rev **00023-spm** + worker on digest **1838c07f** (master **00e8f92**; ships
#531/#535/#537/#538/#539/#529/#534/#548/#547 over 304af7b). **CP DOWNGRADED 4/16 ? 2 vCPU/8
GiB** (safe via #539 acquire-gating; val/seg/compile all on L4). Image-only bump preserves the
baked cloud-serving/phase_placement config. Verified 200 on /api/config. Rollback ? digest
**7b62a8df** @ 4/16 or rev **00022-29w**. Release-process gaps ? **532**.
5      - ? **CLOSED 2026-07-08:** [followup-validation-latency-live-cuj.md](followup-
validation-latency-live-cuj.md) ? live SPA CUJ filled both CF-gated cells on rev **00019**
(validation wall **183 s** PASS + reel compiled; fresh clip, corroborates not re-measures
Video-1). Docs **PR #522**. Surfaced stitch bottleneck ? issues **521** (compile?L4 + config
placement) & **520** (debug tracing).
6
7      - ? **Perf gate + nightly reorg (PR #547 MERGED 2026-07-09, squash `00e8f92f`):** [perf-
release-gate.md](perf-release-gate.md) ? local release perf gate (prod vs HEAD, per-stage GPU
timing, **stitch-only gating**) at `scripts/gates/perf/`; nightly CLIs consolidated into
`scripts/gates/{perf,quality}/`; OS-aware (WSL on Windows, `detector_impl=native` on Linux).
Finding: release **perf-neutral** (#531 within local noise). Related: the pre-existing Windows-
only `test_storage` #535 failure (`task_2e0a3b82`) ? **FIXED in PR #548** ([ci-windows-coverage-
gap.md](ci-windows-coverage-gap.md)).
8
9      - [Actual repo location](actual-repo-location.md) ? primary working dir is an empty
STUB; real code is under `khelsutra-guru/badminton-highlight-indexer`.
10     - [Efficient timeouts, retries & Windows gcloud SSH](efficient-timeouts-and-retries.md)
? no silent retry loops; plain `gcloud compute ssh` (plink/pscp reject `-o`, mishandle `~`);
diagnose-first, run long work detached.
```

122. user (2026-07-09T06:57:37.686Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

123. assistant (2026-07-09T06:57:59.835Z)

Done. PR [#550](https://github.com/Khelsutra/badminton-highlight-indexer/pull/550) is open against `master`, branched from `origin/master @ 00e8f92`.

What was wrong and what I changed

All three gaps are fixed in one PR (`backend/worker/\_\_main\_\_.py`, `backend/pipeline/detectors/native\_wasb\_runner.py`, `backend/pipeline/detectors/wasb\_infer.py`):

****GAP 1 ? validation invisible on PASS.**** The worker now logs `[worker] validating (N checks)?` then `[worker] validation PASSED/FAILED in Ys ? <name>=ok/FAIL(score)?`, and the per-validator

verdicts land in the pushed `report.json` under `validation` (previously only in the ephemeral worker DB).

GAP 2 ? no per-stage timing. checksum / validation / segmentation are each timed with `perf\_counter()`, logged, and persisted as a `timings` block in `report.json` ? built incrementally so a mid-run failure still records the stages that ran.

GAP 3 ? the #534 heartbeat never fired. This turned out to be **two compounding bugs**, which is why the CUJ saw **zero** snapshots:

- **Emission:** `[wasb\_infer.py]` (backend/pipeline/detectors/wasb_infer.py) never configured logging, so its `done/total (pct%)` progress went through `logger.info()` while root stayed at `WARNING` ? the lines were dropped entirely, not just misrouted.
- **Capture:** even when emitted they go to **stderr**, but `[native\_wasb\_runner.py]` (backend/pipeline/detectors/native_wasb_runner.py) streamed **only stdout** (and draining stdout while the child fills its ~64 KB stderr pipe risks a deadlock).

Fix: `wasb\_infer` configures INFO logging to **stdout** at entry, and the runner **merges stderr?stdout** (line-buffered) with a bounded tail for failure diagnostics. The callback now fires ? `run\_telemetry.jsonl` snapshots, the `[worker] progress:` heartbeat, and the `\_detector\_progress.json` sidecar all flow.

`report.json` changes are additive ? the control-plane reads it via `report.get(...)` , so existing readers are unaffected.

Verified locally (Windows)

- `run\_all\_tests.py --cov=backend --cov=training --cov-fail-under=70` ? **1965** passed, **1** skipped, coverage **84.44%** (worker `\_\_main\_\_.py` 91%). No `test\_storage.py` failures ? #548 is already on master.
- `ruff check .` (0.15.16) and `mypy backend training` (2.1.0) ? both clean.
- Added offline tests for the report/timing serialization, the `\_val\_token` score hint, the stderr-merge + failure-tail streaming path, and the `wasb\_infer` stdout logging config.

Not verified (stated plainly)

The **live-L4 heartbeat end-to-end** ? real GPU + `wasb` conda env + Cloud Run L4 ? can't be exercised offline. The unit tests lock the emission config, the stdout parse/merge contract, and the serialization, but confirming the actual telemetry/sidecar/heartbeat on a real pass needs an **owner-driven CUJ**. I noted this in the PR body and in memory.

124. user (2026-07-09T07:12:46.258Z)

<cross-session-message from="local_5780afdd-5e0d-4851-8707-9826eb374e87" name="API performance comparison testing" encoded="1">

Heads-up before you start editing: your cwd `C:\Users\avidu\Projects\badminton-highlight-indexer` is the near-empty STUB (not a git repo). The REAL code is at `C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer` (see the `actual-repo-location` memory). Do all work there.

Concrete anchors for this task in the real repo:

- `backend/worker/\_\_main\_\_.py` ? validation runs ~line 290 (`validator\_registry.run\_all\_with\_context(...)`) but is SILENT on pass and results go only to the ephemeral DB, never report.json. Add per-stage log lines + wall-clock timing (validation / segment / compile) so runlogs can capture "validation took Ns". Also emit a pass summary, not just FAILED/SKIPPED.
- The `\_progress` callback (~line 388, wired `progress=\_progress` ~line 434) writes `&video\_id&detector\_progress.json` + prints `[worker] progress:` ? but on the NATIVE-L4 path it never fires (heartbeat dead: run_telemetry.jsonl was just a 116-byte header on the 07-09 live CUJ). Wire the progress callback through `backend/pipeline/detectors/native\_wasb\_runner.py` so it actually emits on the native path.

Git safety (shared checkout, concurrent sessions): `git -C &real-repo& fetch origin` then `git checkout -b &name& origin/master` ; stage explicit paths (never `git add -A`); before PR, `git log --oneline origin/master..HEAD` must list ONLY your commits. Full local gate before LGTM: run_all_tests.py (--cov-fail-under=70), ruff (pinned 0.15.16) + mypy clean. End commits with the Co-Authored-By: Claude Opus 4.8 trailer.

</cross-session-message>

125. assistant (2026-07-09T07:13:11.926Z)

This heads-up describes the task I've **already completed and shipped** ? it reads as pre-start context, but the work is done. Let me quickly confirm my branch and PR are still intact (concurrent-session sanity check) before responding.

126. user (2026-07-09T07:13:18.368Z)

```
From https://github.com/Khelsutra/badminton-highlight-indexer
  00e8f92..6cfcb37  master    -> origin/master
=== branch ===
fix/worker-l4-infer-observability
=== my commit vs origin/master ===
b3dff78 fix(worker): make the cloud-serving L4 infer path observable
=== PR ===
#550 [OPEN] fix(worker): make the cloud-serving L4 infer path observable (fix/worker-l4-infer-observability)
```

127. assistant (2026-07-09T07:13:26.390Z)

My PR #550 is intact ? but `origin/master` moved (`00e8f92 ? 6cfcb37`) while I worked, so a sibling session merged something. Since the sibling that messaged me was working on the same task, I need to check whether they shipped a duplicate. Let me see what landed.

128. user (2026-07-09T07:13:28.471Z)

```
=== new commits on master ===
6cfcb37 docs(runlog): CP 2vCPU/8Gi downgrade + 7.84 GiB live CUJ (2026-07-09) (#549)

=== files touched by those commits ===
..._cloud-serving_cp-2vcpu-downgrade_2026-07-09.md | 90 +++++
1 file changed, 90 insertions(+)
```

129. assistant (2026-07-09T07:13:55.478Z)

Confirmed ? no duplication or conflict. The new master commit is just [#549](https://github.com/Khelsutra/badminton-highlight-indexer/pull/549), a docs/runlog file, disjoint from my five files, so PR #550 still merges cleanly.

This task is already complete ? the heads-up arrived after I'd shipped it. What the sibling flagged matches exactly what I did, in the real repo (`khelsutra-guru/badminton-highlight-indexer`), branched from `origin/master @ 00e8f92`:

- **GAP 1** (validation silent on pass) ? added `[worker] validating (N checks)?` + a `PASSED/FAILED` summary with per-validator `name=ok/FAIL(score)`, and the verdicts now go into `report.json`, not just the ephemeral DB.
- **GAP 3** (native-L4 heartbeat dead ? the 116-byte header you saw) ? root-caused to two compounding bugs: `wasb\_infer.py` never configured logging (progress `logger.info` suppressed under root=WARNING) **and** it emits to stderr while `native\_wasb\_runner.py` read only stdout. Fixed both (INFO?stdout in `wasb\_infer`, stderr?stdout merge + line-buffering in the runner), so `\_progress` fires ? `run\_telemetry.jsonl` snapshots + `[worker] progress:` + the `\_detector\_progress.jsonl` sidecar.
- Git safety followed (branched from `origin/master`, explicit path staging, only my commit `b3dff78` on the branch); full gate green (**1965** passed, 84.44% cov, ruff 0.15.16 + mypy clean).

One scope nuance worth flagging: your note mentioned per-stage timing for "validation / segment / **compile**". My PR times **checksum / validation / segmentation** in the **infer** path (`cmd\_infer`) ? which is what the task spec listed (GAP 2). The **compile** stage lives in a separate subcommand (`cmd\_compile`, the #521 L4 stitch) and wasn't in this task's scope; its observability is tracked under #520/#521. It's a ~10-line follow-up to add the same `timings`/log pattern to `cmd\_compile` if you want it in this PR ? say the word and I'll fold it

in. Otherwise #550 is ready for review as-is.

130. user (2026-07-09T07:23:07.651Z)

pls review comments on PR

131. assistant (2026-07-09T07:23:17.050Z)

The user wants me to review the comments on PR #550. Let me fetch all of them ? inline review comments, review summaries, and the conversation.

132. user (2026-07-09T07:23:24.505Z)

===== REVIEWS (summaries) =====

--- avidullu [COMMENTED] @ 2026-07-09T07:23:07Z

Review submitted: 1 bug, 3 suggestions, 1 nit. Overall approve-with-notes on the three L4 observability fixes; the main residual risk is the streaming Popen path not enforcing timeout_sec.

===== INLINE REVIEW COMMENTS =====

--- avidullu on backend/pipeline/detectors/native_wasb_runner.py:386 (2026-07-09T07:22:34Z)
[bug] The live-progress branch (entered whenever `progress` is not None) launches the child with `subprocess.Popen` and then blocks on `for line in proc.stdout` / `proc.wait()` with no timeout. `NativeWasbConfig.timeout\_sec` is only applied on the fast path via `\_run()` ? `subprocess.run(..., timeout=...)`. The worker always passes a progress callback (`trajectory\_hybrid` ? `run\_predict(..., progress=progress)`), so the cloud native-L4 path this PR is about never enforces the configured 30?minute kill. The surrounding `except subprocess.TimeoutExpired` is therefore dead code for production streaming, and the new comment?s ?hangs until timeout? framing is misleading. A child that stops producing output (or wedges after closing pipes) can block the worker indefinitely.

Suggestion: Enforce `timeout\_sec` on the streaming path as well ? e.g. track `deadline = time.monotonic() + timeout`, use a reader thread/`select` with a short poll, or `proc.wait(timeout=remaining)` in a loop and `kill()` + raise/return on expiry (mirroring `\_run`). Add a unit test that a fake Popen which never closes stdout is killed when `timeout\_sec` is small.

--- avidullu on backend/worker/__main__.py:359 (2026-07-09T07:22:34Z)

[suggestion] `n\_checks = len(getattr(validator\_registry, "\_validators", []))` reaches into a private registry attribute. That works today, but it is not part of the public validator-registry API and will silently report `0 checks` if the internal field is renamed or if validators are stored differently.

Suggestion: Prefer a public helper (e.g. `len(validator\_registry)` / `list\_registered()` if one exists or is added next to `register`), or log without a count (`[worker] validating?`) to avoid coupling to private state.

--- avidullu on backend/worker/__main__.py:369 (2026-07-09T07:22:34Z)

[suggestion] On validation failure the new summary line (`validation FAILED in ? ? name=FAIL(score)`) is always followed by the pre-existing failures print later in `cmd\_infer` (`validation FAILED: name: message`). Operators get two nearly-duplicate FAIL lines with different formats in one runlog, which slightly muddies the ?loud single verdict? goal.

Suggestion: Either drop/replace the older failures print now that the summary line exists, or make the new line the sole PASS path log and keep the detailed `name: message` form only on FAIL (not both).

--- avidullu on backend/worker/__main__.py:330 (2026-07-09T07:22:34Z)

[suggestion] Stage timings cover checksum / validation / segmentation but not the storage pull (`storage.fetch` of the input video), which is often a large fraction of cloud wall-clock for multi?GB clips. The same mid-run `\_fail()` path therefore still omits the stage that may dominate before checksum starts.

****Suggestion:**** Time the fetch into something like ``timings["fetch_s"]`` (and optionally push) so runlogs explain full end-to-end latency, not only post-download stages.

```
--- avidullu on backend/worker/__main__.py:374 (2026-07-09T07:22:34Z)
**[nit]** On `--skip-validation`, `validation_payload` is `[]` and is still written as
`"validation": []` (because the gate is `is not None`). That is backward-compatible and distinct
from omitting the key, but a later CP/P3 consumer cannot tell ?dispatcher already validated
(skipped here)? from ?worker ran zero validators.?
```

****Suggestion:**** Optionally add ``"validation_skipped": true`` (or omit ``validation`` when skipped) so future control-plane ingest of worker verdicts does not treat an empty list as an authoritative empty suite. Not blocking for this observability-only PR.

===== ISSUE/CONVERSATION COMMENTS =====

133. assistant (2026-07-09T07:27:39.132Z)

The owner (avidullu) left a substantive review: 1 bug, 3 suggestions, 1 nit. Let me address all of them ? the timeout bug is the "main residual risk" and must be fixed. First, a quick check for the ``__len__`` truthiness concern before I add a public registry accessor.

134. user (2026-07-09T07:27:41.647Z)

No matches found

135. assistant (2026-07-09T07:27:49.126Z)

No boolean use of the registry anywhere ? ``__len__`` is safe. Now let me implement all the review fixes.

****Fix 1 (the bug):** enforce ``timeout_sec`` on the streaming path.** Add ``threading``, and use a watchdog Timer that kills the child at the deadline ? the kill closes stdout, the read loop hits EOF, and I raise ``TimeoutExpired`` through the existing handler (making it live for this path, not dead).

136. user (2026-07-09T07:27:51.321Z)

The file `C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\detectors\native_wasb_runner.py` has been updated successfully. (file state is current in your context ? no need to Read it back)

137. user (2026-07-09T07:28:15.033Z)

The file `C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\detectors\native_wasb_runner.py` has been updated successfully. (file state is current in your context ? no need to Read it back)

138. assistant (2026-07-09T07:28:21.101Z)

Let me verify the full streaming block + exception handlers are coherent after the edit:

139. user (2026-07-09T07:28:21.878Z)

```
367
368         res: Optional[subprocess.CompletedProcess] = None
369         try:
370             if progress is None:
371                 # Fast path: capture at end (existing behaviour, byte-identical for
tests).
372                 res = self._run(argv, cwd=self._repo_src())
373             else:
374                 # Live progress path (#534): stream the child's output so we can parse
the
```

```

375 # periodic "done/total (pct%) ..." lines wasb_infer emits and fire the
callback ?
376 # a real heartbeat during the long GPU pass instead of a spinner.
377 #
378 # MERGE stderr INTO stdout (stderr=STDOUT). wasb_infer routes its
logging (progress
379 # included) to a stream, and third-party libs may also write to stderr.
Reading
380 # stdout alone would (a) miss any progress that lands on stderr and (b)
risk a
381 # DEADLOCK: draining only stdout while the child fills its ~64 KB stderr
pipe blocks
382 # the child on the stderr write, which stops its stdout too ? the run
then hangs
383 # until timeout. One merged pipe removes both failure modes (the
2026-07-09 live CUJ
384 # produced zero heartbeat lines). bufsize=1 = line-buffered so the
heartbeat is
385 # timely, not block-delayed. A bounded tail keeps the last lines for the
failure
386 # message, since there is no separate stderr to read on error.
387 proc = subprocess.Popen(
388     argv,
389     cwd=self._repo_src(),
390     stdout=subprocess.PIPE,
391     stderr=subprocess.STDOUT,
392     text=True,
393     bufsize=1,
394 )
395 assert proc.stdout is not None
396 tail: "deque[str]" = deque(maxlen=200)
397 # TIMEOUT: the blocking ``for line in proc.stdout`` read has NO built-in
deadline, so a
398 # wedged child (hung GPU call, or one that closed its pipes without
exiting) would
399 # block the worker indefinitely ? the fast ``_run()`` path's
``timeout_sec`` is not in
400 # force here, and this IS the production path (the worker always passes
a progress
401 # callback). A watchdog Timer kills the child at ``timeout_sec``; the
kill closes its
402 # stdout, the read loop hits EOF, and we surface a timeout via the SAME
403 # ``TimeoutExpired`` handler as ``_run`` (return None, log "timed out
after Ns").
404 timed_out = threading.Event()
405 watchdog: Optional[threading.Timer] = None
406 if self.cfg.timeout_sec:
407
408     def _kill_on_timeout() -> None:
409         timed_out.set()
410         try:
411             proc.kill()
412         except Exception: # noqa: BLE001 - the run is already being
torn down
413             pass
414
415     watchdog = threading.Timer(self.cfg.timeout_sec, _kill_on_timeout)
416     watchdog.daemon = True
417     watchdog.start()
418 try:
419     for line in proc.stdout:
420         line = line.rstrip()
421         if not line:
422             continue
423         tail.append(line)

```

```

424         # Parse wasb_infer progress lines for the detector heartbeat.
425         # Example: "1234/5000 (24%) | 2.3 win/s | elapsed 45s | ETA
120s"
426         m = re.search(r"(\d+)/(\d+)\s+(\d+)\s+", line)
427         if m:
428             done, total, pct = m.groups()
429             try:
430                 progress(f"detector {pct}% ({done}/{total})")
431             except Exception: # noqa: BLE001
432                 pass # progress must never kill the run
433         # Mirror to the container/runlog.
434         logger.info(line)
435     proc.wait()
436     res = subprocess.CompletedProcess(
437         args=argv,
438         returncode=proc.returncode,
439         stdout="",
440         stderr="\n".join(tail),
441     )

```

140. user (2026-07-09T07:28:25.926Z)

```

441         )
442     finally:
443         if watchdog is not None:
444             watchdog.cancel()
445         if proc.poll() is None:
446             proc.kill()
447             proc.wait()
448         if timed_out.is_set():
449             # Route through the outer handler so the streaming path enforces
timeout_sec
450             # exactly like _run (no longer dead code for the production
native-L4 path).
451             raise subprocess.TimeoutExpired(argv, self.cfg.timeout_sec)
452     except FileNotFoundError:
453         logger.error(
454             "python binary not found: %r (set WASB_PYTHON).", self.cfg.python_bin
455         )
456         return None
457     except subprocess.TimeoutExpired:
458         logger.error(
459             "WASB native inference timed out after %ss.", self.cfg.timeout_sec
460         )
461         return None
462
463     if res is None or res.returncode != 0:
464         logger.error(
465             "WASB native inference failed:\n%s",
466             (getattr(res, "stderr", None) or getattr(res, "stdout", None) or
"")[-2000:],
467         )
468         return None
469     if not os.path.exists(out_csv):
470         logger.error(

```

141. assistant (2026-07-09T07:28:35.419Z)

The timeout fix is coherent ? it routes through the existing `TimeoutExpired` handler, making it live for the streaming path.

****Fix 2: public registry count.**** Add `\_\_len\_\_` to `ValidationRegistry`:

142. user (2026-07-09T07:28:38.508Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\validators\base.py has been updated successfully. (file state is current in your context ? no need to Read it back)

143. assistant (2026-07-09T07:28:53.025Z)

****Fix 4:** time the storage pull (``fetch_s``)** ? it can dominate wall-clock for multi-GB clips, and the mid-run ``_fail()`` should capture it. Moving ``timings`` init above the pull:

144. user (2026-07-09T07:29:02.247Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

145. assistant (2026-07-09T07:29:18.210Z)

****Fixes 2, 3, 5 in the validation block:**** public ``len(validator_registry)``, a ``skipped`` flag driving ``validation_skipped`` in the report (instead of an ambiguous ``[]``), and setting up to reword the duplicate FAIL line:

146. user (2026-07-09T07:29:27.378Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

147. assistant (2026-07-09T07:29:34.873Z)

****Fix 3:** reword the duplicate FAIL line** so it complements the verdict summary instead of repeating "validation FAILED":

148. user (2026-07-09T07:29:38.121Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

149. assistant (2026-07-09T07:29:44.498Z)

****Fix 5:** ``validation_skipped`` in the report writer.** Add the param and emit it:

150. user (2026-07-09T07:29:56.241Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

151. assistant (2026-07-09T07:30:05.470Z)

Now thread ``validation_skipped`` through ``_serialize_and_push_outputs``:

152. user (2026-07-09T07:30:16.328Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

153. assistant (2026-07-09T07:30:25.428Z)

Now update the two push sites (the ``_fail`` closure and the success path) to pass ``report_validation`` + ``validation_skipped``:

154. user (2026-07-09T07:30:28.286Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

155. user (2026-07-09T07:30:34.294Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

156. assistant (2026-07-09T07:31:24.945Z)

All source fixes are in. Now let me update and add the tests. Let me check the validators-base test file's style first, then update the worker/native tests.

157. user (2026-07-09T07:31:25.948Z)

```
1      """Tests for the validation registry, focused on the backward-compatible
2      `run_all_with_context` added for observability reporting."""
3
4      from backend.config.models import AppConfig
5      from backend.pipeline.validators.base import (
6          ValidationRegistry,
7          ValidationResult,
8          VideoValidator,
9      )
10
11
12      class _ProducerValidator(VideoValidator):
13          @property
14          def name(self):
15              return "producer"
16
17          @property
18          def description(self):
19              return "writes ffprobe_meta into the shared context"
20
21          def validate(self, video_path, config, context):
22              context["ffprobe_meta"] = {"fps": 59.94, "width": 1920, "height": 1080}
23              return ValidationResult validator_name=self.name, passed=True, message="ok")
24
25
26      class _ConsumerValidator(VideoValidator):
27          @property
28          def name(self):
29              return "consumer"
30
31          @property
32          def description(self):
33              return "reads ffprobe_meta the producer cached"
34
35          def validate(self, video_path, config, context):
36              seen = "ffprobe_meta" in context
37              return ValidationResult(
38                  validator_name=self.name, passed=seen, message=f"saw_meta={seen}"
39              )
40
```

158. user (2026-07-09T07:31:33.870Z)

```
@property
def description(self):
```



```

        return "Tests safe instance registration for validators with complex __init__
state."

    def validate(self, video_path, config, context):
        # Use the state to "prove" it was initialized once
        return ValidationResult(
            validator_name=self.name,
            passed=len(self._heavy_state) > 0,
            message="ok",
        )

    r = ValidationRegistry()
    v = ComplexValidator()
    r.register(v)

    results, _ = r.run_all_with_context("v.mp4", {})
    assert len(results) == 1
    assert results[0].validator_name == "complex-validator"
    assert results[0].passed is True

```

159. user (2026-07-09T07:31:39.542Z)

appended

160. assistant (2026-07-09T07:31:45.170Z)

Now the native-runner timeout test ? appending after the streaming tests:

161. user (2026-07-09T07:31:54.282Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_native_wasb_runner.py has been updated successfully. (file state is current in your context ? no need to Read it back)

162. assistant (2026-07-09T07:32:11.313Z)

Now update the worker tests for `fetch\_s`, the de-duplicated FAIL line, and the `validation\_skipped` marker. First the PASS test's timings set:

163. user (2026-07-09T07:32:14.615Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

164. user (2026-07-09T07:32:23.095Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

165. assistant (2026-07-09T07:32:26.317Z)

Now the `validation\_skipped` assertions ? extend the omit test and the existing skip-bypass test:

166. user (2026-07-09T07:32:28.808Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

167. user (2026-07-09T07:32:35.046Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_worker.py has been updated successfully. (file state is current in your context ? no need to Read it back)

168. assistant (2026-07-09T07:32:40.429Z)

All fixes and tests are in. Let me run the affected test files, then the full gate.

169. user (2026-07-09T07:32:48.605Z)

```
..... [ 66%]
..... [100%]
108 passed in 4.78s
```

170. assistant (2026-07-09T07:32:57.709Z)

108 pass (including the 0.3s timeout test). Now ruff + mypy:

171. user (2026-07-09T07:32:59.444Z)

All checks passed!

172. user (2026-07-09T07:33:03.636Z)

```
backend\annotations\base.py:88: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\sports\__init__.py:9: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\pipeline\validators\base.py:123: note: By default the bodies of untyped functions are
not checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 150 source files
```

173. assistant (2026-07-09T07:33:11.457Z)

Ruff + mypy clean. Now the full suite with coverage (the pre-push gate):

174. user (2026-07-09T07:39:57.340Z)

```
backend\worker\__init__.py          0      0 100%
backend\worker\__main__.py         463    43  91%
training\__init__.py                0      0 100%
training\gen0\__init__.py           0      0 100%
training\gen0\harness.py           130     9  93%
-----
TOTAL                             13607  2116  84%
Required test coverage of 70% reached. Total coverage: 84.45%
1967 passed, 1 skipped, 57 warnings in 401.43s (0:06:41)
=====
  RESULT: PASS - all tests green
=====
```

175. assistant (2026-07-09T07:40:14.658Z)

Full suite: **1967 passed, 1 skipped**, coverage **84.45%**, worker `\_\_main\_\_.py` 91%. ruff + mypy clean. All five review items addressed and verified.

Let me re-verify branch safety, then stage explicit paths and commit the review fixes.

176. user (2026-07-09T07:40:19.736Z)

```
=== branch ===
fix/worker-l4-infer-observability
=== dirty files ===
  M backend/pipeline/detectors/native_wasb_runner.py
  M backend/pipeline/validators/base.py
```

```
M backend/worker/__main__.py
M tests/test_native_wasb_runner.py
M tests/test_validators_base.py
M tests/test_worker.py
?? "SESSION_HANDOFF (1).md"
=== my commits vs origin/master (b3dff78 should be the only one so far) ===
b3dff78 fix(worker): make the cloud-serving L4 infer path observable
```

177. assistant (2026-07-09T07:40:25.184Z)

Clean state. Staging the 6 explicit paths (the stray untracked handoff excluded) and committing:

178. user (2026-07-09T07:40:26.909Z)

warning: in the working copy of 'tests/test_validators_base.py', LF will be replaced by CRLF the next time Git touches it

```
=== staged ===
M backend/pipeline/detectors/native_wasb_runner.py
M backend/pipeline/validators/base.py
M backend/worker/__main__.py
M tests/test_native_wasb_runner.py
M tests/test_validators_base.py
M tests/test_worker.py
?? "SESSION_HANDOFF (1).md"
```

179. assistant (2026-07-09T07:40:31.707Z)

Staged (the LF/CRLF warning is benign normalization). Committing:

180. user (2026-07-09T07:40:43.070Z)

```
[fix/worker-l4-infer-observability e54e4a1] fix(worker): address PR review ? streaming timeout +
observability polish
6 files changed, 153 insertions(+), 19 deletions(-)
```

181. assistant (2026-07-09T07:40:46.972Z)

Committed `e54e4a1`. Verifying my commits and pushing:

182. user (2026-07-09T07:40:50.493Z)

```
=== my commits vs origin/master (both mine) ===
e54e4a1 fix(worker): address PR review ? streaming timeout + observability polish
b3dff78 fix(worker): make the cloud-serving L4 infer path observable
=== push ===
To https://github.com/Khelsutra/badminton-highlight-indexer.git
b3dff78..e54e4a1 fix/worker-l4-infer-observability -> fix/worker-l4-infer-observability
```

183. assistant (2026-07-09T07:40:57.703Z)

Both commits pushed cleanly. Now let me post a reply on the PR mapping each review finding to its resolution:

184. user (2026-07-09T07:41:12.136Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/550#issuecomment-4922778714>

185. assistant (2026-07-09T07:41:22.879Z)

Review response posted. Let me update the memory to reflect the review round for future sessions.

186. user (2026-07-09T07:41:31.872Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\pr550-worker-l4-observability.md has been updated successfully. (file state is current in your context ? no need to Read it back)

187. user (2026-07-09T07:41:40.466Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

188. assistant (2026-07-09T07:41:58.251Z)

Done. I reviewed the owner's review on [PR #550](https://github.com/Khelsutra/badminton-highlight-indexer/pull/550) ? ****1 bug, 3 suggestions, 1 nit**** ? and addressed all five in commit `e54e4a1` (pushed), with a per-item [response comment](https://github.com/Khelsutra/badminton-highlight-indexer/pull/550#issuecomment-4922778714).

The findings and how I resolved each

****? [BUG] streaming path ignored `timeout\_sec`**** ? the real risk. The live-progress branch is the **production** native-L4 path (the worker always passes a callback), yet `for line in proc.stdout` / `wait()`` had no deadline ? a wedged child could hang the worker forever, and the `except subprocess.TimeoutExpired`` was dead code there. Fixed with a watchdog `threading.Timer`` that kills the child at `timeout_sec``; the kill closes stdout ? the read hits EOF ? I raise through the ****same**** `TimeoutExpired`` handler as `_run`` (`return None``, logs "timed out after `Ns"`). New test drives a fake child that blocks until the watchdog kills it.

****[suggestion] private `\_validators` access**** ? added `ValidationRegistry.__len__`` and use `len(validator_registry)`` (verified nothing relies on registry truthiness).

****[suggestion] two near-duplicate FAIL lines**** ? the detail print is reworded to `[worker] validation rejected by: ?``, complementing the `validation FAILED in Ys ? ?`` verdict line; test locks a single-verdict occurrence.

****[suggestion] pull stage untimed**** ? added `timings["fetch_s"]`` (often the dominant term for multi-GB clips; a mid-run `_fail()`` now captures it).

****[nit] ambiguous `validation: []` on skip**** ? skip now writes `validation_skipped: true`` and omits the array, so a consumer can tell "dispatcher already validated" from "zero-validator suite".

Verification

Re-ran the full gate: ****1967 passed / 1 skipped****, coverage ****84.45%**** (floor 70), ruff 0.15.16 + mypy 2.1.0 clean. Two commits on the branch, both mine, on top of `origin/master``.

The one unchanged caveat I flagged in the response: the ****live-L4 heartbeat end-to-end**** still needs an owner-driven CUJ to confirm on real hardware ? everything else is verified offline.

189. user (2026-07-09T07:50:09.883Z)

ptal more comments on the PR

190. assistant (2026-07-09T07:50:14.932Z)

Let me fetch the latest comments and reviews on the PR.

191. user (2026-07-09T07:50:21.784Z)

===== REVIEWS (summaries, chronological) =====

--- avidullu [COMMENTED] @ 2026-07-09T07:23:07Z

Review submitted: 1 bug, 3 suggestions, 1 nit. Overall approve-with-notes on the three L4 observability fixes; the main residual risk is the streaming Popen path not enforcing `timeout_sec``.

--- avidullu [COMMENTED] @ 2026-07-09T07:47:45Z

Second-pass review submitted: all five prior items resolved. One residual suggestion on a narrow timeout-watchdog race (success can lose to a late Timer fire); not a ship-blocker at the default 1800s budget.

===== INLINE REVIEW COMMENTS =====

--- avidullu on backend/pipeline/detectors/native_wasb_runner.py:387 @ 2026-07-09T07:22:34Z (commit b3dff78)

[bug] The live-progress branch (entered whenever `progress` is not None) launches the child with `subprocess.Popen` and then blocks on `for line in proc.stdout` / `proc.wait()` with no timeout. `NativeWasbConfig.timeout\_sec` is only applied on the fast path via `\_run()` ? `subprocess.run(..., timeout=...)`. The worker always passes a progress callback (`trajectory\_hybrid` ? `run\_predict(..., progress=progress)`), so the cloud native-L4 path this PR is about never enforces the configured 30-minute kill. The surrounding `except subprocess.TimeoutExpired` is therefore dead code for production streaming, and the new comment?s ?hangs until timeout? framing is misleading. A child that stops producing output (or wedges after closing pipes) can block the worker indefinitely.

Suggestion: Enforce `timeout\_sec` on the streaming path as well ? e.g. track `deadline = time.monotonic() + timeout`, use a reader thread/`select` with a short poll, or `proc.wait(timeout=remaining)` in a loop and `kill()` + raise/return on expiry (mirroring `\_run`). Add a unit test that a fake Popen which never closes stdout is killed when `timeout\_sec` is small.

--- avidullu on backend/worker/__main__.py:359 @ 2026-07-09T07:22:34Z (commit b3dff78)

[suggestion] `n\_checks = len(getattr(validator\_registry, "\_validators", []))` reaches into a private registry attribute. That works today, but it is not part of the public validator-registry API and will silently report `0 checks` if the internal field is renamed or if validators are stored differently.

Suggestion: Prefer a public helper (e.g. `len(validator\_registry)` / `list\_registered()`) if one exists or is added next to `register`, or log without a count (`[worker] validating?`) to avoid coupling to private state.

--- avidullu on backend/worker/__main__.py:384 @ 2026-07-09T07:22:34Z (commit b3dff78)

[suggestion] On validation failure the new summary line (`validation FAILED in ? ? name=FAIL(score)`) is always followed by the pre-existing failures print later in `cmd\_infer` (`validation FAILED: name: message`). Operators get two nearly-duplicate FAIL lines with different formats in one runlog, which slightly muddies the ?loud single verdict? goal.

Suggestion: Either drop/replace the older failures print now that the summary line exists, or make the new line the sole PASS path log and keep the detailed `name: message` form only on FAIL (not both).

--- avidullu on backend/worker/__main__.py:332 @ 2026-07-09T07:22:34Z (commit b3dff78)

[suggestion] Stage timings cover checksum / validation / segmentation but not the storage pull (`storage.fetch` of the input video), which is often a large fraction of cloud wall-clock for multi-GB clips. The same mid-run `\_fail()` path therefore still omits the stage that may dominate before checksum starts.

Suggestion: Time the fetch into something like `timings["fetch\_s"]` (and optionally push) so runlogs explain full end-to-end latency, not only post-download stages.

--- avidullu on backend/worker/__main__.py:392 @ 2026-07-09T07:22:34Z (commit b3dff78)

[nit] On `--skip-validation`, `validation\_payload` is `[]` and is still written as `"validation": []` (because the gate is `is not None`). That is backward-compatible and distinct from omitting the key, but a later CP/P3 consumer cannot tell ?dispatcher already validated (skipped here)? from ?worker ran zero validators.?

Suggestion: Optionally add `"validation\_skipped": true` (or omit `validation` when skipped) so future control-plane ingest of worker verdicts does not treat an empty list as an authoritative empty suite. Not blocking for this observability-only PR.

--- avidullu on backend/pipeline/detectors/native_wasb_runner.py:409 @ 2026-07-09T07:47:13Z

```
(commit e54e4a1)
**[suggestion]** The watchdog callback always does `timed_out.set()` before `proc.kill()`, and the post-loop check is an unconditional `if timed_out.is_set(): raise TimeoutExpired`. If the child exits successfully (rc=0, CSV written) in the narrow window after `proc.wait()` / before `watchdog.cancel()`, a late timer fire can still set `timed_out` (while `Popen.kill()` is a no-op once `returncode` is set). The run is then reported as a timeout (`return None`) despite a clean completion ? success loses the race to the timer. Cancel-on-completion and daemon Timer are otherwise correct; kill?EOF?raise path works for a truly wedged child; no timer leak when the process dies first.
```

```
**Suggestion:** Only claim timeout when the child was still alive (and/or when completion was not clean), e.g. set the flag only if `proc.poll()` is None` before kill, and/or after the try/finally treat success as winning the race: `if timed_out.is_set() and not (res is not None and res.returncode == 0): raise TimeoutExpired(...)`. Optionally log the merged `tail` on the timeout path (today `res.stderr` is built then discarded by the raise) for post-mortem context.
```

```
===== ISSUE/CONVERSATION COMMENTS =====
```

```
--- avidullu @ 2026-07-09T07:41:13Z
```

```
Thanks for the review ? all five addressed in `e54e4a1`.
```

```
**[bug] streaming path ignored `timeout_sec` (`native_wasb_runner.py`).** Fixed. The live-progress branch now arms a watchdog `threading.Timer(timeout_sec, ?)` that `proc.kill()`s the child at the deadline; the kill closes stdout ? the `for line in proc.stdout` read hits EOF ? I `raise subprocess.TimeoutExpired` so it flows through the same `except` handler as `_run` (`return None`, `"timed out after Ns"`). So that handler is no longer dead code for the production native-L4 path, and a silent/wedged child is killed instead of hanging. New test `test_run_predict_streaming_enforces_timeout` drives a fake child whose stdout read blocks until the watchdog kills it and asserts `out is None` + the kill fired + the timeout was logged.
```

```
**[suggestion] private `_validators` access (`__main__.py:359`).** Fixed. Added a public `ValidationRegistry.__len__` and switched to `len(validator_registry)`. (Confirmed nothing uses the registry in a boolean context, so `__len__` introduces no truthiness surprise.) Test: `test_registry_len_counts_registered_validators`.
```

```
**[suggestion] two near-duplicate FAIL lines (`__main__.py:369`).** Fixed. The verdict line stays `[worker] validation FAILED in Ys ? name=FAIL(score)`; the pre-existing detail print is reworded to `[worker] validation rejected by: name: message` so it complements rather than repeats. The FAIL test now asserts `count("[worker] validation FAILED") == 1` plus the distinct `rejected by:` line.
```

```
**[suggestion] pull stage untimed (`__main__.py:330`).** Fixed. `storage.fetch` is now timed into `timings["fetch_s"]` (and logged on the pull line), so a mid-run `_fail()` explains latency that accrued before checksum ? often the dominant term for multi-GB clips.
```

```
**[nit] ambiguous `validation: []` on skip (`__main__.py:374`).** Fixed. On `--skip-validation` the report now carries `validation_skipped: true` and omits the `validation` array, so a CP/P3 consumer can distinguish "dispatcher already validated" from "worker ran a zero-validator suite". Tests updated for the skip report shape.
```

```
Local gate re-run green: **1967 passed / 1 skipped**, coverage **84.45%** (floor 70), ruff 0.15.16 + mypy 2.1.0 clean. The live-L4 heartbeat end-to-end still needs an owner CUJ (unchanged caveat).
```

192. assistant (2026-07-09T07:52:32.557Z)

Good ? the second pass confirms all five prior items resolved, with **one new, narrow correctness finding**: a watchdog race where a clean success finishing right at the deadline could be misreported as a timeout (the timer sets `timed\_out` after `res` is built, and my post-check raises unconditionally). It's real, so I'll fix it properly ? success must win the race. Let me apply the reviewer's suggested guards.

193. user (2026-07-09T07:52:37.709Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-

indexer\backend\pipeline\detectors\native_wasb_runner.py has been updated successfully. (file state is current in your context ? no need to Read it back)

194. assistant (2026-07-09T07:52:40.990Z)

Now the post-check ? success wins the race, and carry the tail for post-mortem:

195. user (2026-07-09T07:52:45.051Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\detectors\native_wasb_runner.py has been updated successfully. (file state is current in your context ? no need to Read it back)

196. assistant (2026-07-09T07:52:48.385Z)

Now the handler ? log the tail on timeout:

197. user (2026-07-09T07:52:52.031Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\backend\pipeline\detectors\native_wasb_runner.py has been updated successfully. (file state is current in your context ? no need to Read it back)

198. assistant (2026-07-09T07:53:16.376Z)

Now update the timeout test to also verify the tail is surfaced, and add a test for the success-wins-the-race case:

199. user (2026-07-09T07:53:32.338Z)

The file C:\Users\avidu\Projects\khelsutra-guru\badminton-highlight-indexer\tests\test_native_wasb_runner.py has been updated successfully. (file state is current in your context ? no need to Read it back)

200. assistant (2026-07-09T07:53:41.890Z)

Let me run the affected tests plus ruff/mypy:

201. user (2026-07-09T07:53:55.488Z)

```
..... [ 66%]
..... [100%]
109 passed in 4.81s
===== RUFF =====
All checks passed!
===== MYPY =====
backend\annotations\base.py:88: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\sports\__init__.py:9: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\pipeline\validators\base.py:123: note: By default the bodies of untyped functions are
not checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:9: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 150 source files
```

202. assistant (2026-07-09T07:54:03.087Z)

109 passed (+1 race test), ruff + mypy clean. Running the full suite as the pre-push gate:

203. user (2026-07-09T08:00:36.295Z)

training\gen0\harness.py 130 9 93%

```

-----
TOTAL                                13609   2114   84%
Required test coverage of 70% reached. Total coverage: 84.47%
1968 passed, 1 skipped, 57 warnings in 388.81s (0:06:28)
=====
RESULT: PASS - all tests green
=====

```

204. assistant (2026-07-09T08:00:42.502Z)

Full suite green: **1968 passed, 1 skipped**, coverage **84.47%**. Let me self-review the final race-fix diff, then commit.

205. user (2026-07-09T08:00:43.706Z)

```

diff --git a/backend/pipeline/detectors/native_wasb_runner.py
b/backend/pipeline/detectors/native_wasb_runner.py
index 5cbf3a4..d48aea5 100644
--- a/backend/pipeline/detectors/native_wasb_runner.py
+++ b/backend/pipeline/detectors/native_wasb_runner.py
@@ -406,11 +406,15 @@ class NativeWasbRunner(DetectorRunner):
     if self.cfg.timeout_sec:

         def _kill_on_timeout() -> None:
-            timed_out.set()
-            try:
-                proc.kill()
-            except Exception: # noqa: BLE001 - the run is already being torn down
-                pass
+            # Only claim a timeout if the child is STILL running: a late fire in
+            # the
+            # window after a clean exit must not flag `timed_out` (proc.kill() is a
+            # no-op once the process has a returncode), so a success can't be
+            # misreported.
+            if proc.poll() is None:
+                timed_out.set()
+                try:
+                    proc.kill()
+                except Exception: # noqa: BLE001 - the run is already tearing down
+                    pass

         watchdog = threading.Timer(self.cfg.timeout_sec, _kill_on_timeout)
         watchdog.daemon = True
@@ -445,18 +449,29 @@ class NativeWasbRunner(DetectorRunner):
     if proc.poll() is None:
         proc.kill()
         proc.wait()
-    if timed_out.is_set():
-        # Route through the outer handler so the streaming path enforces
-        # exactly like _run (no longer dead code for the production native-L4
-        # path).
-        raise subprocess.TimeoutExpired(argv, self.cfg.timeout_sec)
+    # A CLEAN completion wins the race against a late timer fire: only surface a
+    # timeout
+    # when the run did NOT finish cleanly (rc 0). Route through the outer handler
+    # so the
+    # streaming path enforces timeout_sec exactly like _run (that handler is no
+    # longer
+    # dead code for the production native-L4 path); carry the merged tail for a
+    # post-mortem.
+    if timed_out.is_set() and not (res is not None and res.returncode == 0):
+        raise subprocess.TimeoutExpired(
+            argv, self.cfg.timeout_sec, stderr="\n".join(tail)
+        )

```



```

except FileNotFoundError:
    logger.error(
        "python binary not found: %r (set WASB_PYTHON).", self.cfg.python_bin
    )
    return None
- except subprocess.TimeoutExpired:
+ except subprocess.TimeoutExpired as e:
+     # Surface the merged output tail (the streaming path passes it as `stderr`) so a
timeout
+     # is diagnosable ? not just "timed out after Ns".
+     detail = next(
+         (p for p in (e.stderr, e.output) if isinstance(p, str) and p), ""
+     )
    logger.error(
-         "WASB native inference timed out after %ss.", self.cfg.timeout_sec
+         "WASB native inference timed out after %ss.%s",
+         self.cfg.timeout_sec,
+         f"\nlast output:\n{detail[-2000:]}" if detail else "",
    )
    return None

```

206. assistant (2026-07-09T08:00:49.144Z)

Clean and coherent ? all three parts of the finding addressed. Let me verify branch state, then commit and push.

207. user (2026-07-09T08:00:53.826Z)

```

=== branch ===
fix/worker-l4-infer-observability
=== dirty (mine) ===
M backend/pipeline/detectors/native_wasb_runner.py
M tests/test_native_wasb_runner.py
?? "SESSION_HANDOFF (1).md"
=== my commits vs origin/master ===
e54e4a1 fix(worker): address PR review ? streaming timeout + observability polish
b3dff78 fix(worker): make the cloud-serving L4 infer path observable

```

208. user (2026-07-09T08:01:05.599Z)

```

[fix/worker-l4-infer-observability f7f0c82] fix(detectors): close the timeout-watchdog race
(success must win a late fire)
2 files changed, 87 insertions(+), 13 deletions(-)

```

209. user (2026-07-09T08:01:11.972Z)

```

=== my commits vs origin/master (all mine) ===
f7f0c82 fix(detectors): close the timeout-watchdog race (success must win a late fire)
e54e4a1 fix(worker): address PR review ? streaming timeout + observability polish
b3dff78 fix(worker): make the cloud-serving L4 infer path observable
=== push ===
To https://github.com/Khelsutra/badminton-highlight-indexer.git
e54e4a1..f7f0c82 fix/worker-l4-infer-observability -> fix/worker-l4-infer-observability

```

210. assistant (2026-07-09T08:01:15.188Z)

Pushed (`f7f0c82`). Now posting the response to the race finding:

211. user (2026-07-09T08:01:26.086Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/550#issuecomment-4922925046>

212. assistant (2026-07-09T08:01:31.476Z)

Comment posted. Let me update the memory with the race fix.

213. user (2026-07-09T08:01:39.139Z)

```
<tool_use_error>String to replace not found in file.  
String: `Ys ? ?`). (4) time the storage pull ? `timings["fetch_s"]` (often the dominant term for  
multi-GB  
clips; mid-run `_fail()` now captures it). (5) skip path writes `validation_skipped: true`  
(omits the  
`validation` array) instead of an ambiguous `validation: []`.  
(note: Edit also tried swapping \uxxxx escapes and their characters; neither form matched, so  
the mismatch is likely elsewhere in old_string. Re-read the file and copy the exact surrounding  
text.)</tool_use_error>
```

214. assistant (2026-07-09T08:01:42.228Z)

Let me re-read the end of the memory file to match exactly.

215. user (2026-07-09T08:01:43.523Z)

```
36      ? 1967 passed / 1 skipped, coverage 84.45% (worker `__main__.py` 91%); ruff 0.15.16 +  
mypy 2.1.0 clean.  
37      No `test_storage.py` failures ? #548 (already on master) fixed those on Windows.  
38  
39      **Review round addressed (commit `e54e4a1`)** ? owner review (1 bug, 3 suggestions, 1  
nit) all fixed:  
40      (1) **BUG:** the streaming (progress) path ? the PRODUCTION native-L4 path, worker  
always passes a  
41      callback ? enforced NO `timeout_sec`; a wedged child could hang forever. Now a watchdog  
42      `threading.Timer` kills the child at the deadline ? EOF ? raise through the same  
`TimeoutExpired`  
43      handler as `_run` (that handler was dead code for streaming before). (2)  
`len(validator_registry)`  
44      via a new `ValidationRegistry.__len__` (no more private `_validators`). (3) de-duped  
FAIL logs ? the  
45      detail line is now `[worker] validation rejected by: ?` (verdict line keeps `validation  
FAILED in  
46      Ys ? ?`). (4) time the storage pull ? `timings["fetch_s"]` (often the dominant term for  
multi-GB  
47      clips; mid-run `_fail()` now captures it). (5) skip path writes `validation_skipped:  
true` (omits the
```

216. user (2026-07-09T08:01:46.861Z)

```
47      clips; mid-run `_fail()` now captures it). (5) skip path writes `validation_skipped:  
true` (omits the  
48      `validation` array) instead of an ambiguous `validation: []`.  
49  
50      **OPEN follow-up (not verifiable offline):** the live-L4 heartbeat end-to-end (real GPU  
+ `wasb` env +  
51      Cloud Run L4) needs an **owner-driven CUJ** to confirm the telemetry snapshots /  
progress sidecar /  
52      stdout heartbeat actually land. The offline tests lock the emission config, the stdout  
parse+merge
```

217. user (2026-07-09T08:01:55.224Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\pr550-worker-l4-observability.md has been updated successfully. (file state is current in your context ? no need to Read it back)

218. user (2026-07-09T08:02:05.677Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-

indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

219. assistant (2026-07-09T08:02:21.690Z)

The second pass is handled. Here's the state of PR
[#550](https://github.com/Khelsutra/badminton-highlight-indexer/pull/550):

Second-pass review outcome

The owner confirmed **all five prior items resolved**, with one new **narrow correctness finding** on the timeout watchdog I'd just added ? a genuine race:

The race: a clean child completion (`rc=0`, CSV written) that finished in the window *before* `watchdog.cancel()` could still catch a late Timer fire. The timer set `timed_out`, and my unconditional `if timed_out.is_set(): raise` then reported that **success as a timeout**. The reviewer correctly noted the `kill()` was the harmless part (no-op once the child has a returncode) ? the *flag* was the bug.

Fix (`f7f0c82`)

Two guards so a clean completion deterministically wins:

- Watchdog only flags a live child** ? `if proc.poll() is None: timed_out.set(); proc.kill()`, so a late fire after a clean exit never sets the flag.
- Success wins the post-check** ? `if timed_out.is_set() and not (res is not None and res.returncode == 0): raise`, so an `rc=0` run is returned as success regardless.

Plus the post-mortem suggestion: the merged output `tail` now rides into `TimeoutExpired(stderr=?)` and gets logged (`"?timed out after Ns. last output: ?"`) instead of being built-then-discarded.

Tests: new `test_run_predict_streaming_clean_exit_beats_late_timer` (forces `timed_out` set + `rc=0` ? asserts success survives); the wedged-child timeout test now also asserts the tail is surfaced. A truly wedged child still returns `None`.

Verification

Full suite **1968 passed / 1 skipped**, coverage **84.47%** (floor 70), ruff 0.15.16 + mypy 2.1.0 clean. Pushed and [replied on the PR](https://github.com/Khelsutra/badminton-highlight-indexer/pull/550#issuecomment-4922925046).

Both review passes were **COMMENT-only** (not blocking `CHANGES_REQUESTED`), so the branch is in good shape ? the remaining gate is the owner's merge decision plus the live-L4 heartbeat CUJ that can only be confirmed on real hardware.

220. user (2026-07-09T08:02:23.878Z)

pls merge the PR after responding to the reviewer and ensuring all tests and CI pass

221. assistant (2026-07-09T08:02:42.722Z)

Merging is a hard-to-reverse, outward-facing action, so let me verify thoroughly first: check for any new reviewer comments, the formal review decision, and ? critically ? the CI status and mergeability.

222. user (2026-07-09T08:02:55.322Z)

```
===== PR STATE =====
#550 state=OPEN draft=false
mergeable=MERGEABLE mergeState=UNSTABLE reviewDecision=
head=fix/worker-l4-infer-observability base=master commits=3
```

```
===== CI CHECKS =====
Built wheel smoke (install + launch)    pass    1m17s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/29003359679/job/86068967570
Golden fixtures cross-OS (macos-latest) pass    22s      https://github.com/Khelsutra/badminton-
```

```
highlight-indexer/actions/runs/29003359679/job/86068967725
Leak scan (PII / attribution guard)      pass    7s      https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967582
Markdown link check (live docs) pass    8s      https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967535
Mypy (lenient baseline) pass    31s      https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967609
Syntax sweep + import smoke (no GPU stack)      pass    34s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967564
Full test suite (offline, mocked)      pending 0      https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967543
Dependency license scan (no copleft in runtime deps)      pass    36s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967639
Golden fixtures cross-OS (ubuntu-latest)      pass    25s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967572
Golden fixtures cross-OS (windows-latest)      pass    56s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967610
Ruff lint      pass    20s      https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967538
```

===== NEW REVIEWS (since last) =====

```
--- avidullu [COMMENTED] @ 2026-07-09T07:23:07Z
--- avidullu [COMMENTED] @ 2026-07-09T07:47:45Z
```

===== LATEST CONVERSATION COMMENTS (authors + time) =====

```
--- avidullu @ 2026-07-09T07:41:13Z
--- avidullu @ 2026-07-09T08:01:27Z
```

===== LATEST INLINE COMMENTS (after f7f0c82) =====

```
--- avidullu on backend/pipeline/detectors/native_wasb_runner.py:387 @ 2026-07-09T07:22:34Z
--- avidullu on backend/worker/__main__.py:359 @ 2026-07-09T07:22:34Z
--- avidullu on backend/worker/__main__.py:384 @ 2026-07-09T07:22:34Z
--- avidullu on backend/worker/__main__.py:332 @ 2026-07-09T07:22:34Z
--- avidullu on backend/worker/__main__.py:392 @ 2026-07-09T07:22:34Z
--- avidullu on backend/pipeline/detectors/native_wasb_runner.py:409 @ 2026-07-09T07:47:13Z
```

223. assistant (2026-07-09T08:03:45.373Z)

PR is `MERGEABLE` but `mergeState=UNSTABLE` because the `***Full test suite (offline, mocked)***` CI check is still `PENDING`? all 10 other checks pass. I've already responded to both review passes (no new reviewer comments since). I must not merge while CI is pending, so let me wait for that check to finish.

224. user (2026-07-09T08:05:09.803Z)

```
Full test suite (offline, mocked)      pass    3m33s      https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967543
Golden fixtures cross-OS (macos-latest) pass    22s      https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967725
Golden fixtures cross-OS (ubuntu-latest)      pass    25s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967572
Golden fixtures cross-OS (windows-latest)      pass    56s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967610
Leak scan (PII / attribution guard)      pass    7s      https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967582
Markdown link check (live docs) pass    8s      https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967535
Mypy (lenient baseline) pass    31s      https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967609
Ruff lint      pass    20s      https://github.com/Khelsutra/badminton-highlight-
```

```

indexer/actions/runs/29003359679/job/86068967538
Syntax sweep + import smoke (no GPU stack)      pass    34s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967564
Built wheel smoke (install + launch)      pass    1m17s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967570
Dependency license scan (no copyleft in runtime deps)      pass    36s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967639
Full test suite (offline, mocked)      pass    3m33s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967543
Golden fixtures cross-OS (macos-latest)      pass    22s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967725
Golden fixtures cross-OS (ubuntu-latest)      pass    25s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967572
Golden fixtures cross-OS (windows-latest)      pass    56s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967610
Leak scan (PII / attribution guard)      pass    7s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967582
Markdown link check (live docs)      pass    8s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/29003359679/job/86068967535
Mypy (lenient baseline)      pass    31s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967609
Ruff lint      pass    20s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967538
Syntax sweep + import smoke (no GPU stack)      pass    34s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/29003359679/job/86068967564
=== exit: 0 ===

```

225. assistant (2026-07-09T08:06:10.817Z)

All 11 CI checks pass (Full test suite: 3m33s). Let me re-confirm the merge state flipped to CLEAN now that CI is green, and double-check the base hasn't moved.

226. user (2026-07-09T08:06:18.556Z)

```

mergeable=MERGEABLE mergeState=CLEAN reviewDecision= state=OPEN
=== base master head ===
6cfcb37
=== does the repo auto-delete merged branches? checking merge_commit settings via API ===
squash_merge_allowed=true delete_branch_on_merge=true

```

227. assistant (2026-07-09T08:06:36.707Z)

Merge state is **CLEAN**, `mergeable`, no blocking review decision, squash allowed, and `delete_branch_on_merge=true` (GitHub will auto-delete the remote branch ? so I'll skip `--delete-branch` to avoid disturbing this shared local checkout). Let me post a brief closing comment, then squash-merge (matching the repo's `(#NNN)` squash convention).

228. user (2026-07-09T08:06:43.225Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/550#issuecomment-4922966425>

229. assistant (2026-07-09T08:07:05.777Z)

Closing comment posted. Now squash-merging with a full commit message (co-author trailer included, matching the repo's `(#NNN)` convention). Not passing `--delete-branch` ? the repo auto-deletes the remote branch, and I want to leave this shared local checkout untouched.

230. user (2026-07-09T08:07:17.718Z)

(Bash completed with no output)

231. assistant (2026-07-09T08:07:20.893Z)

Let me verify the merge landed on master.

232. user (2026-07-09T08:07:28.761Z)

=== PR state ===

Unknown JSON field: "merged"

Available fields:

- additions
- assignees
- author
- autoMergeRequest
- baseRefName
- baseRefOid
- body
- changedFiles
- closed
- closedAt
- closingIssuesReferences
- comments
- commits
- createdAt
- deletions
- files
- fullDatabaseId
- headRefName
- headRefOid
- headRepository
- headRepositoryOwner
- id
- isCrossRepository
- isDraft
- labels
- latestReviews
- maintainerCanModify
- mergeCommit
- mergeStateStatus
- mergeable
- mergedAt
- mergedBy
- milestone
- number
- potentialMergeCommit
- projectCards
- projectItems
- reactionGroups
- reviewDecision
- reviewRequests
- reviews
- state
- statusCheckRollup
- title
- updatedAt
- url

=== fetch + master HEAD ===

From https://github.com/Khelsutra/badminton-highlight-indexer

6cfcb37..01d1081 master -> origin/master

01d1081 fix(worker): make the cloud-serving L4 infer path observable (#550)

6cfcb37 docs(runlog): CP 2vCPU/8Gi downgrade + 7.84 GiB live CUJ (2026-07-09) (#549)

00e8f92 feat(gates): local release perf gate + consolidate nightly into scripts/gates/ (#547)

=== remote branch still there? ===

(empty above = remote branch auto-deleted)

233. assistant (2026-07-09T08:07:38.810Z)

Merge succeeded ? `origin/master` is now `01d1081` = "fix(worker): make the cloud-serving L4 infer path observable (#550)", and the remote branch auto-deleted. Let me confirm the PR shows MERGED and the squash landed exactly my files (no sibling work absorbed).

234. user (2026-07-09T08:07:42.593Z)

```
=== PR state ===
state=MERGED mergedAt=2026-07-09T08:07:18Z by=avidullu sha=01d108112
=== files in the squash commit (must be ONLY my 7) ===
01d1081 fix(worker): make the cloud-serving L4 infer path observable (#550)
 backend/pipeline/detectors/native_wasb_runner.py | 104 ++++++---
 backend/pipeline/detectors/wasb_infer.py          |  22 +++
 backend/pipeline/validators/base.py               |   5 +
 backend/worker/__main__.py                        | 164 ++++++++-----
 tests/test_native_wasb_runner.py                  | 223 ++++++++-----
 tests/test_validators_base.py                     |  10 +
 tests/test_worker.py                              | 186 ++++++++-----
7 files changed, 663 insertions(+), 51 deletions(-)
```

235. assistant (2026-07-09T08:07:52.843Z)

Merge confirmed and clean ? PR **MERGED** (squash `01d1081`), exactly my 7 files, no sibling work absorbed. Let me update the memory to reflect the merged state.

236. user (2026-07-09T08:07:53.956Z)

```
1      ---
2      name: pr550-worker-l4-observability
3      description: "PR #550 (open) makes the cloud-serving L4 worker infer path observable;
live-L4 heartbeat still needs an owner CUJ"
4      metadata:
5        node_type: memory
6        type: project
7        originSessionId: e663947f-2127-4383-9905-650340ff9863
8      ---
9
10     # PR #550 ? cloud-serving L4 infer observability (OPEN, 2026-07-09)
11
12     Branch `fix/worker-l4-infer-observability` cut from `origin/master @ 00e8f92`; commit
`b3dff78`.
```

237. user (2026-07-09T08:07:58.499Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\pr550-worker-l4-observability.md has been updated successfully. (file state is current in your context ? no need to Read it back)

238. user (2026-07-09T08:08:06.310Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\pr550-worker-l4-observability.md has been updated successfully. (file state is current in your context ? no need to Read it back)

239. user (2026-07-09T08:08:17.310Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

240. assistant (2026-07-09T08:08:33.167Z)

Merged. Here's the final state of PR [#550](https://github.com/Khelsutra/badminton-highlight-

indexer/pull/550):

? Merged to master

- **Squash commit** `01d1081` ? "fix(worker): make the cloud-serving L4 infer path observable (#550)".
- Merged by avidullu at 2026-07-09T08:07Z. Remote branch auto-deleted (repo setting); I left the shared local checkout untouched.
- The squash touched **exactly the 7 files** of this PR ? no sibling work absorbed.

What I verified before merging

- **Reviewer feedback**: both COMMENT-only review passes were fully resolved (5 items + the watchdog-race follow-up); no new reviewer comments outstanding. Posted a closing comment summarizing the resolution.
- **CI**: waited for the pending check, then confirmed **all 11 checks green** ? Full test suite (offline/mock, 3m33s), cross-OS golden fixtures (ubuntu/macos/windows), ruff, mypy, wheel smoke, syntax/import sweep, leak + license scans.
- **Local gate**: 1968 passed / 1 skipped, coverage 84.47% (floor 70), ruff 0.15.16 + mypy 2.1.0 clean.
- **Merge state**: `MERGEABLE` / `CLEAN`, no blocking review decision.

?? One thing to flag ? merged ? deployed

The code is on master but **not live yet**. The serving worker/control-plane still run the pre-#550 image (digest `1838c07f` = master `00e8f92`). To make the new observability actually take effect ? and to run the **live-L4 heartbeat CUJ** that's the one thing CI can't cover (real GPU + `wasb` env + Cloud Run L4) ? the worker image needs a rebuild + redeploy. I've recorded this in memory as the open follow-up. Want me to prep that deploy, or leave it for the next release cycle?