

codex-windows / 019f218f-8ff0-7f50-9e90-d8405ca40d12

Metadata

```
- Source: `codex-windows`  
- Kind: `codex`  
- Source file: `C:\Users\avidu\codex\archived_sessions\rollout-2026-07-02T12-11-24-019f218f-8ff0-7f50-9e90-d8405ca40d12.jsonl`  
- SHA-256: `39d58843bfa4091f10ccf37592ae411faec1e3c8961b15bd65a7d021adfab169`  
- Source modified: `2026-07-02T15:51:48+00:00`  
- Imported at: `2026-07-05T16:48:16+00:00`  
- cli_version: `0.142.5`  
- cwd: `C:\Users\avidu\Projects\khelsutra-guru`  
- model_provider: `openai`  
- session_id: `019f218f-8ff0-7f50-9e90-d8405ca40d12`  
- source: `vscode`
```

Transcript

1. developer (2026-07-02T06:41:27.526Z)

```
<permissions instructions>  
Filesystem sandboxing defines which files can be read or written. `sandbox_mode` is `danger-  
full-access`: No filesystem sandboxing - all commands are permitted. Network access is enabled.  
Approval policy is currently never. Do not provide the `sandbox_permissions` for any reason,  
commands will be rejected.  
</permissions instructions>  
<app-context>
```

Codex desktop context

- You are running inside the Codex (desktop) app, which allows some additional features not available in the CLI alone:

Images/Visuals/Files

- In the app, the model can display images and videos using standard Markdown image syntax: `![alt](url)`
- When sending or referencing a local image or video, always use an absolute filesystem path in the Markdown image tag (e.g., `![alt](/absolute/path.png)`); relative paths and plain text will not render the media.
- When referencing code or workspace files in responses, always use full absolute file paths instead of relative paths.
- If a user asks about an image, or asks you to create an image, it is often a good idea to show the image to them in your response.
- Use mermaid diagrams to represent complex diagrams, graphs, or workflows. Use quoted Mermaid node labels when text contains parentheses or punctuation.
- Return web URLs as Markdown links (e.g., `[label](https://example.com)`).

Workspace Dependencies

- For sheets, slides, and documents, call ``load_workspace_dependencies`` to find the bundled runtime and libraries.

Automations

- This app supports recurring automations, reminders, monitors, follow-ups, and thread wakeups. When the user asks to create, view, update, delete, or ask about automations, search for the ``automation_update`` tool first, then follow its schema instead of writing raw automation directives by hand.

- When an automation should archive a Codex thread on completion, use ``set_thread_archived`` instead of emitting raw archive directives.

Thread Coordination

- When the user asks to create, fork, inspect, continue, hand off, pin, archive, rename, or otherwise manage Codex threads, search for the relevant thread tool first: ``create_thread``, ``fork_thread``, ``list_threads``, ``read_thread``, ``send_message_to_thread``, ``handoff_thread``,

`set\_thread\_pinned`, `set\_thread\_archived`, or `set\_thread\_title`.

- Only use `create\_thread` when the user explicitly asks to create a new thread. Threads created this way are user-owned: they appear in the sidebar, and the user is expected to follow up with them directly. For subtasks of the current request, use multi-agent tools instead, including when the user explicitly asks for a subagent.
- After a successful `create\_thread` call, emit `::created-thread{threadId="..."}` for a created thread or `::created-thread{pendingWorktreeId="..."}` for queued worktree setup on its own line in your final response.

Inline Code Comments

- Use the `::code-comment{...}` directive when you need to attach feedback directly to specific code lines.
- Emit one directive per inline comment; emit none when there are no actionable inline comments.
- Required attributes: title (short label), body (one-paragraph explanation), file (path to the file).
- Optional attributes: start, end (1-based line numbers), priority (0-3).
- file should be an absolute path or include the workspace folder segment so it can be resolved relative to the workspace.
- Keep line ranges tight; end defaults to start.
- Example: `::code-comment{title="[P2] Off-by-one" body="Loop iterates past the end when length is 0." file="/path/to/foo.ts" start=10 end=11 priority=2}`

```
</app-context>
<collaboration_mode># Collaboration Mode: Default
```

You are now in Default mode. Any previous instructions for other modes (e.g. Plan mode) are no longer active.

Your active mode changes only when new developer instructions with a different ``<collaboration_mode>...</collaboration_mode>`` change it; user requests or tool descriptions do not change mode by themselves. Known mode names are Default and Plan.

request_user_input availability

Use the `request\_user\_input` tool only when it is listed in the available tools for this turn.

In Default mode, strongly prefer making reasonable assumptions and executing the user's request rather than stopping to ask questions. If you absolutely must ask a question because the answer cannot be discovered from local context and a reasonable assumption would be risky, ask the user directly with a concise plain-text question. Never write a multiple choice question as a textual assistant message.

```
</collaboration_mode>
<apps_instructions>
```

Apps (Connectors)

Apps (Connectors) can be explicitly triggered in user messages in the format `[\${app-name}](app://{connector\_id})`. Apps can also be implicitly triggered as long as the context suggests usage of available apps.

An app is equivalent to a set of MCP tools within the `codex\_apps` MCP.

An installed app's MCP tools are either provided to you already, or can be lazy-loaded through the `tool\_search` tool. If `tool\_search` is available, the apps that are searchable by `tools\_search` will be listed by it.

Do not additionally call `list\_mcp\_resources` or `list\_mcp\_resource\_templates` for apps.

```
</apps_instructions>
<skills_instructions>
```

Skills

A skill is a set of local instructions to follow that is stored in a `SKILL.md` file. Below is the list of skills that can be used. Each entry includes a name, description, and a short path that can be expanded into an absolute path using the skill roots table.

Skill roots

- `r0` = `C:/Users/avidu/.agents/skills`
- `r1` = `C:/Users/avidu/.codex/skills/.system`
- `r2` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/anthropic-skills/1.0.0/skills`
- `r3` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/cowork-plugin-management/0.2.2/skills`

- `r4` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/design/1.2.0/skills`
- `r5` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/engineering/1.2.0/skills`
- `r6` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/productivity/1.3.0/skills`
- `r7` = `C:/Users/avidu/.codex/plugins/cache/openai-bundled`
- `r8` = `C:/Users/avidu/.codex/plugins/cache/openai-curated-remote/github/0.1.5/skills`
- `r9` = `C:/Users/avidu/.codex/plugins/cache/openai-curated-remote/gmail/0.1.3/skills`
- `r10` = `C:/Users/avidu/.codex/plugins/cache/openai-curated-remote/google-calendar/1.2.3/skills`
- `r11` = `C:/Users/avidu/.codex/plugins/cache/openai-curated-remote/google-drive/0.1.7/skills`
- `r12` = `C:/Users/avidu/.codex/plugins/cache/openai-primary-runtime`

Available skills

- imagegen: Generate or edit raster images when the task benefits from AI-created bitmap visuals such as photos, illustrations, textures, sprites, mockups, or transparent-background cutouts. Use when Codex should create a brand-new image, transform an existing image, or derive visual variants from references, and the output should be a bitmap asset rather than repo-native code or vector. Do not use when the task is better handled by editing e (file: r1/imagegen/SKILL.md)
- openai-docs: Use when the user asks how to build with OpenAI products or APIs, asks about Codex itself or choosing Codex surfaces, needs up-to-date official documentation with citations, help choosing the latest model for a use case, or model upgrade and prompt-upgrade guidance; use OpenAI docs MCP tools for non-Codex docs questions, use the Codex manual helper first for broad Codex self-knowledge, and restrict fallback browsing to official (file: r1/openai-docs/SKILL.md)
- plugin-creator: Create and scaffold plugin directories for Codex with a required `.codex-plugin/plugin.json`, optional plugin folders/files, valid manifest defaults, and personal-marketplace entries by default. Use when Codex needs to create a new personal plugin, add optional plugin structure, generate or update marketplace entries for plugin ordering and availability metadata, or update an existing local plugin during development with the C (file: r1/plugin-creator/SKILL.md)
- skill-creator: Guide for creating effective skills. This skill should be used when users want to create a new skill (or update an existing skill) that extends Codex's capabilities with specialized knowledge, workflows, or tool integrations. (file: r1/skill-creator/SKILL.md)
- skill-installer: Install Codex skills into \$CODEX_HOME/skills from a curated list or a GitHub repo path. Use when a user asks to list installable skills, install a curated skill, or install a skill from another repo (including private repos). (file: r1/skill-installer/SKILL.md)
- anthropic-skills:consolidate-memory: Reflective pass over your memory files ? merge duplicates, fix stale facts, prune the index. (file: r2/consolidate-memory/SKILL.md)
- anthropic-skills:doc-coauthoring: Guide users through a structured workflow for co-authoring documentation. Use when user wants to write documentation, proposals, technical specs, decision docs, or similar structured content. This workflow helps users efficiently transfer context, refine content through iteration, and verify the doc works for readers. Trigger when user mentions writing docs, creating proposals, drafting specs, or similar documentation tasks. (file: r2/doc-coauthoring/SKILL.md)
- anthropic-skills:docx: Use this skill whenever the user wants to create, read, edit, or manipulate Word documents (.docx files). Triggers include: any mention of 'Word doc', 'word document', '.docx', or requests to produce professional documents with formatting like tables of contents, headings, page numbers, or letterheads. Also use when extracting or reorganizing content from .docx files, inserting or replacing images in documents, performing find-an (file: r2/docx/SKILL.md)
- anthropic-skills:new-session-hello: I want claude to read the github repos linked and have some context when I start a new coding session (file: r2/new-session-hello/SKILL.md)
- anthropic-skills:pdf: Use this skill whenever the user wants to do anything with PDF files. This includes reading or extracting text/tables from PDFs, combining or merging multiple PDFs into one, splitting PDFs apart, rotating pages, adding watermarks, creating new PDFs, filling PDF forms, encrypting/decrypting PDFs, extracting images, and OCR on scanned PDFs to make them searchable. If the user mentions a .pdf file or asks to produce one, use this (file: r2/pdf/SKILL.md)
- anthropic-skills:pptx: Use this skill any time a .pptx file is involved in any way ? as input, output, or both. This includes: creating slide decks, pitch decks, or presentations; reading, parsing, or extracting text from any .pptx file (even if the extracted content will be used elsewhere, like in an email or summary); editing, modifying, or updating existing presentations; combining or splitting slide files; working with templates, layouts, speaker (file: r2/pptx/SKILL.md)
- anthropic-skills:schedule: Create or update a scheduled task that runs automatically. Use when the user says things like "every day", "each morning", "remind me in an hour", "run this at noon", or wants to reschedule an existing task. (file: r2/schedule/SKILL.md)

- anthropic-skills:setup-cowork: Guided Cowork setup ? install role-matched plugins, connect your tools, try a skill. (file: r2/setup-cowork/SKILL.md)
- anthropic-skills:skill-creator: Create new skills, modify and improve existing skills, and measure skill performance. Use when users want to create a skill from scratch, edit, or optimize an existing skill, run evals to test a skill, benchmark skill performance with variance analysis, or optimize a skill's description for better triggering accuracy. (file: r2/skill-creator/SKILL.md)
- anthropic-skills:xlsx: Use this skill any time a spreadsheet file is the primary input or output. This means any task where the user wants to: open, read, edit, or fix an existing .xlsx, .xlsm, .csv, or .tsv file (e.g., adding columns, computing formulas, formatting, charting, cleaning messy data); create a new spreadsheet from scratch or from other data sources; or convert between tabular file formats. Trigger especially when the user references a spr (file: r2/xlsx/SKILL.md)
- browser:control-in-app-browser: Control the in-app Browser. Use to open, navigate, inspect, test, click, type, screenshot, or verify local targets such as localhost, 127.0.0.1, ::1, file://, the current in-app browser tab, and websites shown side by side inside Codex. (file: r7/browser/26.623.81905/skills/control-in-app-browser/SKILL.md)
- chrome:control-chrome: Control the user's Chrome browser for tasks that depend on existing Chrome state: tabs, logged-in sessions, or extensions. Prefer purpose-built connectors, APIs, or CLIs when available. (file: r7/chrome/26.623.81905/skills/control-chrome/SKILL.md)
- computer-use:computer-use: Control Windows apps from Codex (file: r7/computer-use/26.623.81905/skills/computer-use/SKILL.md)
- cowork-plugin-management:cowork-plugin-customizer: Customize a Claude Code plugin for a specific organization's tools and workflows. Use when: customize plugin, set up plugin, configure plugin, tailor plugin, adjust plugin settings, customize plugin connectors, customize plugin skill, tweak plugin, modify plugin configuration. (file: r3/cowork-plugin-customizer/SKILL.md)
- cowork-plugin-management:create-cowork-plugin: Guide users through creating a new plugin from scratch in a cowork session. Use when users want to create a plugin, build a plugin, make a new plugin, develop a plugin, scaffold a plugin, start a plugin from scratch, or design a plugin. This skill requires Cowork mode with access to the outputs directory for delivering the final .plugin file. (file: r3/create-cowork-plugin/SKILL.md)
- design:accessibility-review: Run a WCAG 2.1 AA accessibility audit on a design or page. Trigger with "audit accessibility", "check a11y", "is this accessible?", or when reviewing a design for color contrast, keyboard navigation, touch target size, or screen reader behavior before handoff. (file: r4/accessibility-review/SKILL.md)
- design:design-critique: Get structured design feedback on usability, hierarchy, and consistency. Trigger with "review this design", "critique this mockup", "what do you think of this screen?", or when sharing a Figma link or screenshot for feedback at any stage from exploration to final polish. (file: r4/design-critique/SKILL.md)
- design:design-handoff: Generate developer handoff specs from a design. Use when a design is ready for engineering and needs a spec sheet covering layout, design tokens, component props, interaction states, responsive breakpoints, edge cases, and animation details. (file: r4/design-handoff/SKILL.md)
- design:design-system: Audit, document, or extend your design system. Use when checking for naming inconsistencies or hardcoded values across components, writing documentation for a component's variants, states, and accessibility notes, or designing a new pattern that fits the existing system. (file: r4/design-system/SKILL.md)
- design:research-synthesis: Synthesize user research into themes, insights, and recommendations. Use when you have interview transcripts, survey results, usability test notes, support tickets, or NPS responses that need to be distilled into patterns, user segments, and prioritized next steps. (file: r4/research-synthesis/SKILL.md)
- design:user-research: Plan, conduct, and synthesize user research. Trigger with "user research plan", "interview guide", "usability test", "survey design", "research questions", or when the user needs help with any aspect of understanding their users through research. (file: r4/user-research/SKILL.md)
- design:ux-copy: Write or review UX copy ? microcopy, error messages, empty states, CTAs. Trigger with "write copy for", "what should this button say?", "review this error message", or when naming a CTA, wording a confirmation dialog, filling an empty state, or writing onboarding text. (file: r4/ux-copy/SKILL.md)
- documents:documents: Create, edit, redline, and comment on .docx, Word, and Google Docs-targeted document artifacts inside the container, with a strict render-and-verify workflow. Use `render\_docx.py` to generate page PNGs (and optional PDF) for visual QA, then iterate until layout is flawless before delivering the final document. (file: r12/documents/26.630.12135/skills/documents/SKILL.md)

- engineering:architecture: Create or evaluate an architecture decision record (ADR). Use when choosing between technologies (e.g., Kafka vs SQS), documenting a design decision with trade-offs and consequences, reviewing a system design proposal, or designing a new component from requirements and constraints. (file: r5/architecture/SKILL.md)
- engineering:code-review: Review code changes for security, performance, and correctness. Trigger with a PR URL or diff, "review this before I merge", "is this code safe?", or when checking a change for N+1 queries, injection risks, missing edge cases, or error handling gaps. (file: r5/code-review/SKILL.md)
- engineering:debug: Structured debugging session ? reproduce, isolate, diagnose, and fix. Trigger with an error message or stack trace, "this works in staging but not prod", "something broke after the deploy", or when behavior diverges from expected and the cause isn't obvious. (file: r5/debug/SKILL.md)
- engineering:deploy-checklist: Pre-deployment verification checklist. Use when about to ship a release, deploying a change with database migrations or feature flags, verifying CI status and approvals before going to production, or documenting rollback triggers ahead of time. (file: r5/deploy-checklist/SKILL.md)
- engineering:documentation: Write and maintain technical documentation. Trigger with "write docs for", "document this", "create a README", "write a runbook", "onboarding guide", or when the user needs help with any form of technical writing ? API docs, architecture docs, or operational runbooks. (file: r5/documentation/SKILL.md)
- engineering:incident-response: Run an incident response workflow ? triage, communicate, and write postmortem. Trigger with "we have an incident", "production is down", an alert that needs severity assessment, a status update mid-incident, or when writing a blameless postmortem after resolution. (file: r5/incident-response/SKILL.md)
- engineering:standup: Generate a standup update from recent activity. Use when preparing for daily standup, summarizing yesterday's commits and PRs and ticket moves, formatting work into yesterday/today/blockers, or structuring a few rough notes into a shareable update. (file: r5/standup/SKILL.md)
- engineering:system-design: Design systems, services, and architectures. Trigger with "design a system for", "how should we architect", "system design for", "what's the right architecture for", or when the user needs help with API design, data modeling, or service boundaries. (file: r5/system-design/SKILL.md)
- engineering:tech-debt: Identify, categorize, and prioritize technical debt. Trigger with "tech debt", "technical debt audit", "what should we refactor", "code health", or when the user asks about code quality, refactoring priorities, or maintenance backlog. (file: r5/tech-debt/SKILL.md)
- engineering:testing-strategy: Design test strategies and test plans. Trigger with "how should we test", "test strategy for", "write tests for", "test plan", "what tests do we need", or when the user needs help with testing approaches, coverage, or test architecture. (file: r5/testing-strategy/SKILL.md)
- github:gh-address-comments: Address actionable GitHub pull request review feedback. Use when the user wants to inspect unresolved review threads, requested changes, or inline review comments on a PR, then implement selected fixes. Use the GitHub app for PR metadata and flat comment reads, and use the bundled GraphQL script via `gh` whenever thread-level state, resolution status, or inline review context matters. (file: r8/gh-address-comments/SKILL.md)
- github:gh-fix-ci: Use when a user asks to debug or fix failing GitHub PR checks that run in GitHub Actions. Use the GitHub app from this plugin for PR metadata and patch context, and use `gh` for Actions check and log inspection before implementing any approved fix. (file: r8/gh-fix-ci/SKILL.md)
- github:github: Triage and orient GitHub repository, pull request, and issue work through the connected GitHub app. Use when the user asks for general GitHub help, wants PR or issue summaries, or needs repository context before choosing a more specific GitHub workflow. (file: r8/github/SKILL.md)
- github:yeet: Publish local changes to GitHub by confirming scope, committing intentionally, pushing the branch, and opening a draft PR through the GitHub app from this plugin, with `gh` used only as a fallback where connector coverage is insufficient. (file: r8/yeet/SKILL.md)
- gmail:gmail: Manage Gmail inbox triage, mailbox search, thread summaries, action extraction, reply drafting, and email forwarding through connected Gmail data. Use when the user wants to inspect a mailbox or thread, search email with Gmail query syntax, summarize messages, extract decisions and follow-ups, prepare replies or forwarded messages, or organize messages with explicit confirmation before send, archive, delete, or label actions. (file: r9/gmail/SKILL.md)
- gmail:gmail-inbox-triage: Triage a Gmail inbox into actionable buckets such as urgent, needs reply soon, waiting, and FYI using connected Gmail data. Use when the user asks to triage the inbox, rank what needs attention, find what still needs a reply, or separate important mail from noise. (file: r9/gmail-inbox-triage/SKILL.md)

- google-calendar:google-calendar: Manage scheduling and conflicts in connected Google Calendar data. Use when the user wants to inspect calendars, compare availability, review conflicts, find a meeting room, review event notes or attachments, add or adjust reminders, place temporary holds, or draft exact create, update, reschedule, or cancel changes with timezone-aware details. (file: r10/google-calendar/SKILL.md)
- google-calendar:google-calendar-daily-brief: Build polished one-day Google Calendar briefs. Use when the user asks for today, tomorrow, or a specific date summary with an agenda, conflict flags, free windows, remaining-meeting readouts, or a calendar brief, and the Google Calendar connector is available. (file: r10/google-calendar-daily-brief/SKILL.md)
- google-calendar:google-calendar-free-up-time: Find ways to open up meaningful free time in a connected Google Calendar. Use when the user wants to clear up their day, make room for focus time, create a longer uninterrupted block, or see the smallest set of calendar changes that would give time back. (file: r10/google-calendar-free-up-time/SKILL.md)
- google-calendar:google-calendar-group-scheduler: Find and rank good meeting times for multiple people using connected Google Calendar data. Use when the user wants to schedule a group meeting, compare candidate slots across several attendees, find the best compromise time, or add a room check after narrowing the attendee-compatible options. (file: r10/google-calendar-group-scheduler/SKILL.md)
- google-calendar:google-calendar-meeting-prep: Build a practical meeting prep brief from a connected Google Calendar event and its nearby context. Use when the user wants to prepare for an upcoming meeting, understand what to read beforehand, pull in linked notes or docs, or get a concise brief on what the meeting appears to require. (file: r10/google-calendar-meeting-prep/SKILL.md)
- google-drive:google-docs: Connector-first Google Docs creation and editing in local Codex plugin sessions, with direct native create and batchUpdate workflows for simple docs, DOCX-first import for polished deliverables, target-document checks, smart chip and building-block reconstruction, connector-readback verification, and reference routing for formatting, citations, tables, and write-safety. (file: r11/google-docs/SKILL.md)
- google-drive:google-drive: Use connected Google Drive as the single entrypoint for Drive, Docs, Sheets, and Slides work. Use when the user wants to find, fetch, organize, share, export, copy, or delete Drive files, or summarize and edit Google Docs, Google Sheets, and Google Slides through one unified Google Drive plugin. (file: r11/google-drive/SKILL.md)
- google-drive:google-drive-comments: Write, reply to, and resolve Google Drive comments on Docs, Sheets, Slides, and Drive files with evidence-backed location context. Use when the user asks to leave comments, review a file with comments, respond to comment threads, or resolve Drive comments. (file: r11/google-drive-comments/SKILL.md)
- google-drive:google-sheets: Analyze and edit connected Google Sheets with range precision. Use when the user wants to create Google Sheets, find a spreadsheet, inspect tabs or ranges, search rows, plan formulas, create or repair charts, clean or restructure tables, write concise summaries, or make explicit cell-range updates. (file: r11/google-sheets/SKILL.md)
- google-drive:google-slides: Google Slides work for finding, reading, summarizing, creating, importing, template following, visual cleanup, source-deck adaptation, structural repair, and content edits in native Slides decks. (file: r11/google-slides/SKILL.md)
- pdf:pdf: Read, create, inspect, render, and verify PDF files where visual layout matters. Use Poppler rendering plus Python tools such as reportlab, pdfplumber, and pypdf for generation and extraction. (file: r12/pdf/26.630.12135/skills/pdf/SKILL.md)
- presentations:Presentations: Create or edit PowerPoint or Google Slides decks (file: r12/presentations/26.630.12135/skills/presentations/SKILL.md)
- productivity:memory-management: Two-tier memory system that makes Claude a true workplace collaborator. Decodes shorthand, acronyms, nicknames, and internal language so Claude understands requests like a colleague would. CLAUDE.md for working memory, memory/ directory for the full knowledge base. (file: r6/memory-management/SKILL.md)
- productivity:start: Initialize the productivity system and open the dashboard. Use when setting up the plugin for the first time, bootstrapping working memory from your existing task list, or decoding the shorthand (nicknames, acronyms, project codenames) you use in your todos. (file: r6/start/SKILL.md)
- productivity:task-management: Simple task management using a shared TASKS.md file. Reference this when the user asks about their tasks, wants to add/complete tasks, or needs help tracking commitments. (file: r6/task-management/SKILL.md)
- productivity:update: Sync tasks and refresh memory from your current activity. Use when pulling new assignments from your project tracker into TASKS.md, triaging stale or overdue tasks, filling memory gaps for unknown people or projects, or running a comprehensive scan to catch todos buried in chat and email. (file: r6/update/SKILL.md)
- skill-creator: Create new skills, modify and improve existing skills. Use when users want to create a skill from scratch, edit, or optimize an existing skill. (file: r0/skill-

creator/SKILL.md)

- spreadsheets:Spreadsheets: Use this skill when a user requests to create, modify, analyze, visualize, or work with spreadsheet files (`.xlsx`, `.xls`, `.csv`, `.tsv`) or Google Sheets-targeted spreadsheet artifacts with formulas, formatting, charts, tables, and recalculation. (file: r12/spreadsheets/26.630.12135/skills/spreadsheets/SKILL.md)
- template-creator:template-creator: Create or update a reusable personal Codex artifact-template skill. Use when the user invokes \$template-creator or asks in natural language to create a template using, from, or based on an attached Word document, PowerPoint presentation, or Excel workbook, or explicitly asks to edit or update a passed artifact-template skill. Do not use for one-off artifact creation from an existing template. (file: r12/template-creator/26.630.12135/skills/template-creator/SKILL.md)

How to use skills

- Discovery: The list above is the skills available in this session (name + description + short path). Skill bodies live on disk at the listed paths after expanding the matching alias from `### Skill roots`.
- Trigger rules: If the user names a skill (with `\$SkillName` or plain text) OR the task clearly matches a skill's description shown above, you must use that skill for that turn. Multiple mentions mean use them all. Do not carry skills across turns unless re-mentioned.
- Missing/blocked: If a named skill isn't in the list or the path can't be read, say so briefly and continue with the best fallback.
- How to use a skill (progressive disclosure):
 - 1) After deciding to use a skill, the main agent must expand the listed short `path` with the matching alias from `### Skill roots`, then open and read its `SKILL.md` completely before taking task actions. If a read is truncated or paginated, continue until EOF.
 - 2) When `SKILL.md` references relative paths (e.g., `scripts/foo.py`), resolve them relative to the directory containing that expanded `SKILL.md` first, and only consider other paths if needed.
 - 3) If `SKILL.md` points to extra folders such as `references/`, use its routing instructions to identify the files required for the task. The main agent must read each required instruction or reference file itself before acting on it. Do not delegate reading, summarizing, or interpreting skill instructions to a subagent. Subagents may still perform task work when the selected skill allows it.
 - 4) If `scripts/` exist, prefer running or patching them instead of retyping large code blocks.
 - 5) If `assets/` or templates exist, reuse them instead of recreating from scratch.
- Coordination and sequencing:
 - If multiple skills apply, choose the minimal set that covers the request and state the order you'll use them.
 - Announce which skill(s) you're using and why (one short line). If you skip an obvious skill, say why.
- Context hygiene:
 - Progressive disclosure applies to selecting relevant files, not partially reading a selected instruction file. Do not load unrelated references, scripts, or assets.
 - Avoid deep reference-chasing: prefer opening only files directly linked from `SKILL.md` unless you're blocked.
 - When variants exist (frameworks, providers, domains), pick only the relevant reference file(s) and note that choice.
- Safety and fallback: If a skill can't be applied cleanly (missing files, unclear instructions), state the issue, pick the next-best approach, and continue.

</skills_instructions>

<plugins_instructions>

Plugins

A plugin is a local bundle of skills, MCP servers, and apps.

How to use plugins

- Skill naming: If a plugin contributes skills, those skill entries are prefixed with `plugin\_name:` in the Skills list.
- MCP naming: Plugin-provided MCP tools keep standard MCP identifiers such as `mcp\_server\_tool`; use tool provenance to tell which plugin they come from.
- Trigger rules: If the user explicitly names a plugin, prefer capabilities associated with that plugin for that turn.
- Relationship to capabilities: Plugins are not invoked directly. Use their underlying skills, MCP tools, and app tools to help solve the task.
- Relevance: Determine what a plugin can help with from explicit user mention or from the plugin-associated skills, MCP tools, and apps exposed elsewhere in this turn.
- Missing/blocked: If the user requests a plugin that does not have relevant callable

capabilities for the task, say so briefly and continue with the best fallback.
</plugins_instructions>

2. user (2026-07-02T06:41:27.526Z)

AGENTS.md instructions

<INSTRUCTIONS>

Global instructions (all projects)

Repo-freshness convention: git pull BEFORE reading

****Always `git pull` a repo before starting to read it**** ? at session start for the primary working directory, and again for any sibling/local repo the moment work touches it.

- ****Why:**** the owner closes every session on a "clean slate", but multiple parallel slates (sessions/machines) exist ? PRs get merged and master moves between sessions, so the local checkout can silently lag the remote truth. Pull first so ramp-up reads (handoff, docs, code) reflect where the remote actually is.
- ****How (safe form):**** `git pull --ff-only` on the current branch (a plain fetch+ff). Never auto-merge, rebase, or discard local state to make a pull succeed.
- ****If it can't fast-forward**** (diverged branch, dirty tree in the way): do NOT force anything ? report the state to the owner and proceed read-only on what's local, flagging that it may be stale. Uncommitted changes are often a sibling session's or the owner's WIP; never clobber them.

Git branching safety: concurrent sessions / shared checkouts

Multiple sessions ? and spawned/background tasks ? may share ONE working checkout and create branches + commit ****concurrently****, so `git branch` / `git log` can change UNDER you between two commands. (Spawned "fresh worktree" tasks are NOT always isolated.) This has caused a real incident: a sibling commit landed in the gap between a branch-point check and `git checkout -b`, so the new branch rode on top of it and a ****squash-merge** silently absorbed the sibling's work**. Protocol ? do this EVERY time, not only when you suspect a collision:

- ****Branch from the REMOTE, explicitly:**** `git fetch origin` then `git checkout -b <name> origin/<base>` (e.g. `origin/master`). Never assume the current branch is the intended base.
- ****Re-verify right before `git add`/commit:**** `git branch --show-current` + `git log --oneline -1` (HEAD is yours). Before opening a PR, `git log --oneline origin/<base>..HEAD` must list ONLY your commits ? if a sibling commit appears, re-cut the branch from `origin/<base>`.
- ****Stage explicit paths**** (`git add <files>`) ? never `git add -A`/`.`/`-u`, so sibling or untracked files can't ride into your commit.
- ****Serialize repo writes:**** don't start a repo-writing spawned/background task while you hold uncommitted work ? commit or push first. If one may be running, treat the tree + current branch as able to shift, and put this branching-safety reminder into the spawned task's own prompt.

Session-handoff convention

Maintain a living ****session handoff**** for each project and use it to resume context quickly across windows.

- ****Where it lives:****
 - If the project has an auto-memory dir (`~/Codex/projects/<project>/memory/`), keep the handoff as `session-handoff.md` there, and list it under a ****"? Start here**** entry at the top of that project's `MEMORY.md`.
 - Otherwise, keep a `SESSION\_HANDOFF.md` at the repo root.
- ****What it contains**** (keep it to ~1 page ? it POINTS to detailed artifacts, never duplicates them):
 1. ****You are here**** ? current state.
 2. ****Next steps / open threads.****
 3. ****Ramp-up kit**** ? the exact files/docs to read to get current.

4. **Key decisions** ? so they aren't relitigated.
- **At session start:** if a handoff exists, read it before starting substantive work, and use it (plus its ramp-up kit) to get up to speed instead of replaying past conversation.
- **Updating it:** keep it current **automatically** ? refresh at meaningful checkpoints (a PR is pushed/merged, a key decision is made, work shifts phase, or a session is wrapping up), so a restart can ramp up without losing context. Do **not** rewrite it every turn ? only when something worth resuming from changed. The phrases **"update the handoff"** / **"refresh the handoff"** force an immediate full rewrite on demand.
- **Keep it honest:** it reflects real, shipped state ? never aspirational. (See each project's attribution/working-agreement note where present.)

Code review ? pre-LGTM gate (applies to every PR/code review, every project)

Do **not** post **LGTM** or otherwise sign off on a pull request until I have personally **verified** the following ? never trust the PR description's claims, reproduce them locally:

1. **Full test suite is green.** Run the whole suite (e.g. `run_all_tests.py` / `pytest``), not just the changed files.
2. **Install any dependency needed** so the suite actually collects and runs ? do not let modules skip or error for missing deps. (If a check genuinely can't run locally ? GPU/WSL/real-video/etc. ? say so explicitly and note what was and wasn't verified.)
3. **Lint + type checks are clean** ? `ruff check .`` and `mypy ...``, matching the repo's CI. Install `ruff`/`mypy`` if absent.
4. **Coverage stays up** ? run with coverage and confirm it stays **at or above the repo's floor** (e.g. the CI `--cov-fail-under`` value) and does **not regress** versus the base branch.
5. **Review findings are resolved** ? bugs and repo-convention/discipline issues addressed.

LGTM is a **verified** stamp, never a rubber stamp. State plainly what was and wasn't verifiable locally.

</INSTRUCTIONS>

<environment_context>

```
<cwd>C:\Users\avidu\Projects\khelsutra-guru</cwd>
<shell>powershell</shell>
<current_date>2026-07-02</current_date>
<timezone>Asia/Calcutta</timezone>
<filesystem><workspace_roots><root>C:\Users\avidu\Projects\khelsutra-
guru</root></workspace_roots><permission_profile type="disabled"><file_system
type="unrestricted" /></permission_profile></filesystem>
</environment_context>
```

3. user (2026-07-02T06:41:27.556Z)

Continue the khelsutra-guru health remediation work.

First, obey the repo freshness rule: open `C:\Users\avidu\Projects\khelsutra-guru`, then for every sub-repo you touch run `git pull --ff-only`` before reading. Also read `C:\Users\avidu\OneDrive\Documents\claude-sync\memory\khelsutra-guru\session-handoff.md` and `current-state.md`.

The docs-only bootstrap PR is merged:

- Engine PR #423, commit e517475
- Live tracker: badminton-highlight-indexer/docs/KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN.md
- Start from §7 of that plan.

Current task: execute the plan wave-wise, one small PR per §7 row. Begin with Wave 1 in badminton-highlight-indexer:

1. P2: fix #340 idempotent processing / Gemini re-bill risk.
2. Then P3 source caching, P4 ffmpeg timeouts, P5 scratch/retention, P6 real gcs.py mypy fix.

Important constraints:

- Branch from remote explicitly: `git fetch origin`` then `git checkout -b <branch> origin/master``.

- Re-verify branch + HEAD before staging/commit.
- Stage explicit paths only.
- Keep coverage at/above repo floor; run repo-appropriate lint/type/tests and record verification in PR.
- If touching khelsutra SPA/UI or UI modules, create and validate screenshot/screencast evidence using the existing khelsutra-screencast convention.
- sports-data-collector has a pre-existing dirty `.gitignore` hunk; do not clobber it.`
- R1-27 is still open despite PR #422: ``python -m mypy backend training --no-incremental`` still fails at `backend/storage/gcs.py:31` when `google-cloud-storage` is installed.

4. assistant (2026-07-02T06:41:39.691Z)

I'll start by honoring the freshness convention exactly: fast-forward the primary checkout first, then read the handoff/state files, and only then inspect the relevant sub-repo after pulling it too.

5. assistant (2026-07-02T06:41:53.416Z)

The workspace folder itself is just a container, not a Git checkout, so there was nothing to fast-forward at that level. I'm pulling the engine sub-repo before reading its plan, while also loading the external handoff notes.

6. assistant (2026-07-02T06:42:11.477Z)

The handoff matches your brief: #423 is merged and \$7 row-by-row is the active lane. I'm checking the engine checkout state next, then I'll cut the first branch directly from ``origin/master`` before reading the implementation surface.

7. assistant (2026-07-02T06:42:27.005Z)

I found one small wrinkle: the live plan has the ``gcs.py`` mypy fix as ``P1``, while your note calls it ``P6``. I'll follow your requested order by starting with ``P2`` now, and I'll keep that numbering mismatch visible when I update the tracker.

8. assistant (2026-07-02T06:42:43.995Z)

The P2 branch is now freshly based on ``origin/master``. I'm going to pull the issue text and trace the ``/api/process`` path, especially job creation, video identity, and any existing `?force?` semantics.

9. assistant (2026-07-02T06:42:57.760Z)

The issues confirm the precise shape: reuse must happen by content hash, default path must not call the segmenter, and an explicit force path must preserve the current re-ingest replacement behavior. I'm narrowing the edit surface to request schema, job persistence/reconstruction, and ``_execute_processing``.

10. assistant (2026-07-02T06:43:09.011Z)

The existing async worker reconstructs request fields from the jobs row, so the force flag can't live only in the Pydantic model. I'm checking whether to make it a schema v10 column or use the existing ``jobs.params`` JSON without disturbing old rows.

11. user (2026-07-02T06:43:10.034Z)

please make sure the test coverages don't regress and keep going and merging the PRs

12. assistant (2026-07-02T06:43:24.629Z)

Yep, that's the bar I'll use: no coverage backslide, full local verification before sign-off, then publish and merge only after the PR checks are green. I'm still on P2 and will keep each row isolated as its own branch/PR.

13. assistant (2026-07-02T06:43:34.094Z)

A small implementation choice: I can use the existing ``jobs.params`` JSON for the new force bit, which avoids a schema bump and keeps old queued jobs compatible. That's looking like the least invasive path, provided tests lock the worker reconstruction boundary.

14. assistant (2026-07-02T06:43:49.637Z)

I'm going to edit three places now: add the request flag, persist/reconstruct it through the async job row, and short-circuit ``_execute_processing`` only when the existing video is processed and has stored play segments. Force will bypass that and keep today's replace-after-confirm semantics.

15. assistant (2026-07-02T06:44:11.808Z)

One subtlety: compute provenance already proves 'cloud' vs 'local.' For a cache hit, I'll stamp ``path: cached`` plus ``requested``, so the job result honestly says nothing ran instead of pretending it was local.

16. assistant (2026-07-02T06:45:08.034Z)

I'm applying the code/test changes now. The tests will exercise the real public ``/api/process`` boundary with the inline executor so we catch both the no-force cache hit and the forced reprocess path through the same async machinery the SPA uses.

17. assistant (2026-07-02T06:45:52.705Z)

I'm choosing the stronger cache behavior: after resolving and hashing the input, a processed video with stored segments returns immediately before validators or segmenters. The content hash is the proof; ``force=true`` is the escape hatch when someone deliberately wants to re-run with new settings.

18. assistant (2026-07-02T06:46:27.840Z)

The force bit is now persisted using ``jobs.params`` only when true, so normal process jobs keep the old NULL shape. I'm adding the worker reconstruction and regression tests next.

19. assistant (2026-07-02T06:46:58.316Z)

The no-force test will deliberately make validators and the segmenter 'explode if touched'; that gives us a clean proof that the paid path is not merely hidden behind a success response.

20. assistant (2026-07-02T06:47:40.580Z)

The first test patch caught an accidental insertion point where I'd split an existing test's trailing assertions; I've cleaned that up before running. Focused tests next, then I'll broaden to lint/type/full suite with coverage.

21. assistant (2026-07-02T06:47:56.862Z)

Focused tests show the new default cache behavior correctly changes older 're-ingest safety' tests: those tests were intentionally re-running the same content without an explicit opt-in. I'm updating them to use ``force=true`` so they still cover the failure/crash safety path deliberately.

22. assistant (2026-07-02T06:48:21.278Z)

The async job suite is green now. I'm running the neighboring coverage around job reconstruction, cloud compute provenance, and database schema next because those are the surfaces this touched.

23. assistant (2026-07-02T06:48:45.725Z)

The related suites are also green. I'm doing a quick polish pass on the new params parsing so it remains tolerant if a future caller hands the worker already-decoded params.

24. assistant (2026-07-02T06:48:54.394Z)

That robustness tweak made the line a bit dense; I'm checking formatting/lint next and will let the repo's tools shape it.

25. assistant (2026-07-02T06:49:51.256Z)

The edited surfaces are green again. I'm checking the repo's own CI commands now so the final verification matches the coverage floor and type/lint gates instead of guessing.

26. assistant (2026-07-02T06:50:08.458Z)

The repo's CI floor is `--cov-fail-under=70`; I'm running that full covered suite now, alongside whole-repo ruff and the CI mypy command. If mypy trips the known optional-GCS `gcs.py` issue, I'll separate that clearly from this P2 change.

27. assistant (2026-07-02T06:52:40.127Z)

Full coverage run reached 85.16% but failed two legacy tests for the same reason as before: they were deliberately trying to re-run an already-processed hash without the new force opt-in. Ruff is green; mypy hit the known `backend/storage/gcs.py:31` optional-GCS issue, separate from P2. I'm updating those tests to make their reprocess intent explicit.

28. assistant (2026-07-02T06:52:52.349Z)

Those endpoint tests now state their intent with `force=true`, so the default no-cost reuse path and the deliberate reprocess path are no longer fighting each other. I'm rerunning the endpoint file before another full covered pass.

29. assistant (2026-07-02T06:53:23.615Z)

Endpoint tests are green. I'm rerunning the full covered suite now; coverage was already above floor on the failing run, but I want the clean pass number before publishing.

30. assistant (2026-07-02T06:55:44.215Z)

Full covered suite is now green: 1551 passed, 11 skipped, total coverage 85.16% against the 70% floor. I'm rerunning whole-repo ruff after the last test edits, then I'll prepare the commit and PR.

31. assistant (2026-07-02T06:55:53.685Z)

Local working tree is scoped to the five expected files. Before I commit/publish, I'm loading the GitHub PR workflow skill so the push/PR/merge steps follow the repo's connector conventions rather than freelancing.

32. assistant (2026-07-02T06:56:02.942Z)

I've loaded the publish workflow. Since the branch already follows your explicit base-from-remote rule, I'm staying on it, reviewing the scoped diff, then I'll stage only those five paths.

33. assistant (2026-07-02T06:56:07.932Z)

Scope is clean and intentional: just API/job persistence plus endpoint tests. I'm checking for the GitHub app PR tool; if it isn't exposed cleanly here, I'll use `gh` as the skill's fallback.

34. assistant (2026-07-02T06:56:16.634Z)

The GitHub connector can create the PR once the branch is pushed. I'm at the commit boundary

now, so I'm doing the required branch/HEAD re-check, then I'll stage only the five changed files.

35. assistant (2026-07-02T06:56:20.403Z)

Branch is `fix/process-idempotency` at `e517475`, the docs-only plan commit, so the branch base is clean. Staging the scoped files explicitly now.

36. assistant (2026-07-02T06:56:29.512Z)

Commit is made. I'm verifying that the branch contains only this commit relative to `origin/master` before pushing.

37. assistant (2026-07-02T06:56:37.869Z)

`origin/master..HEAD` is exactly one commit, as required. I'm pushing the branch, then I'll open the PR against `master`.

38. assistant (2026-07-02T06:56:51.027Z)

The branch is on GitHub. I'm creating the PR with the local verification spelled out, including the known local mypy/GCS failure so there's no mystery in the check record.

39. assistant (2026-07-02T06:57:04.471Z)

The connector hit a 404 on PR creation despite the branch existing, so I'm falling back to `gh pr create` as the GitHub publish skill allows.

40. assistant (2026-07-02T06:57:28.299Z)

PR #424 is open. I'm adding the tracker update into that same PR now so S7 reflects P0 as merged and P2 as the row this PR completes.

41. assistant (2026-07-02T06:57:42.551Z)

The tracker change is docs-only. I'm running the markdown link check/diff check, then I'll amend it into the existing P2 commit so the PR stays tidy.

42. assistant (2026-07-02T06:57:46.316Z)

Docs links are clean; branch/HEAD still point at my P2 commit. I'm amending the tracker update into that commit and force-pushing the PR branch.

43. assistant (2026-07-02T06:58:04.121Z)

The PR branch is updated with a single amended commit. I'm refreshing the PR verification text for the docs link check, then I'll watch CI until it's mergeable.

44. assistant (2026-07-02T06:58:10.463Z)

PR #424 is updated and still has exactly one commit over `origin/master`. I'm checking the GitHub runs now; if anything fails, I'll fix it before merging.

45. assistant (2026-07-02T07:01:02.289Z)

All GitHub checks on #424 are green, including the full covered suite and CI mypy. I'm doing one last PR-state check, then I'll merge it.

46. assistant (2026-07-02T07:01:09.139Z)

#424 is clean and not draft. Merging now with the branch deletion so master can move forward for P3.

47. assistant (2026-07-02T07:01:29.728Z)

P2 is merged as #424 and the local checkout fast-forwarded to the new master commit. I'm checking the branch state now, then I'll refresh the handoff before cutting the next row.

48. assistant (2026-07-02T07:01:37.693Z)

Master is now `9523d17`, and #330/#340 are tied to the merged PR. I'm refreshing the handoff/current-state memory before starting P3 so the next window lands in the right place.

49. assistant (2026-07-02T07:01:59.619Z)

The handoff patch missed because the wording in the file differs slightly from the remembered line. I'm reading the small top section and will patch against the exact text.

50. assistant (2026-07-02T07:02:30.728Z)

The persistent handoff is updated. I'm adding the matching short current-state checkpoint now, with P3 as the next concrete row.

51. assistant (2026-07-02T07:02:53.518Z)

Memory is refreshed. Starting P3 now: I'm fast-forwarding the engine repo again, then I'll cut a new branch directly from `origin/master` for the remote source cache work.

52. assistant (2026-07-02T07:03:07.355Z)

P3 branch is fresh off the merged master. I'm tracing `/api/source`, annotation proxy warming, and the storage acquisition API so the cache lands in the narrowest place.

53. assistant (2026-07-02T07:03:17.013Z)

The root cause is as the audit described: source serving falls through to `acquire\_local\_input`, which creates a fresh temp download. I'm reading the exact source/proxy endpoints and their tests now so the cache works for both direct source and proxy generation.

54. assistant (2026-07-02T07:04:29.553Z)

I'm adding a per-video source cache under the configured output base. Remote sources will fetch once into that cache, serve that stable local file for future Range requests, and the proxy warmer will use that cached local copy too.

55. assistant (2026-07-02T07:04:59.233Z)

The code path is in place. I'm adding an API test that simulates two browser source requests against a `gs://` video and proves only one backend fetch happens, while the proxy warmer receives the cached local copy.

56. assistant (2026-07-02T07:05:25.327Z)

The regression test is in. It uses inline threads for determinism and two Range-style requests, so it proves the cache survives the exact browser-scrub shape without waiting on real background work.

57. assistant (2026-07-02T07:05:36.917Z)

The annotation/source tests are green. I'm checking the broader API/storage-adjacent tests before the full covered suite, mostly to catch any interaction with compile/process remote input paths.

58. assistant (2026-07-02T07:06:07.917Z)

The broader API/storage-adjacent tests are green too. The diff is intentionally small: one

source-cache helper path plus one regression test. I'm running the full covered suite and whole-repo lint now.

59. assistant (2026-07-02T07:09:25.648Z)

I'm resuming the P3 source-cache branch. I'll first re-establish the exact branch/tree state and read the GitHub publish workflow skill, then I'll finish verification with coverage in mind before staging anything.

60. assistant (2026-07-02T07:09:33.164Z)

I found the branch exactly where expected: `fix/source-cache` at `9523d17` with only the two P3 files modified so far. I'm going to run the freshness pull on this sub-repo now, then inspect the diff and finish the publish path.

61. assistant (2026-07-02T07:09:43.681Z)

Freshness check passed; the feature branch is still current against `origin/master`. Next I'm reading the current diff and the tracker row so the PR stays small and the docs reflect the row it resolves.

62. assistant (2026-07-02T07:09:57.849Z)

The P3 diff is tightly scoped: one source-cache implementation and one regression test proving only one remote fetch across repeated Range loads. I'm doing the type check now too; based on the handoff I expect the known `backend/storage/gcs.py:31` failure locally, but I'll record it honestly either way.

63. assistant (2026-07-02T07:10:34.855Z)

The local mypy result matches the known R1-27 blocker and points only at `backend/storage/gcs.py:31`. That's outside this P3 row, so I'm keeping this PR focused and will list the failure as pre-existing in the PR verification.

64. assistant (2026-07-02T07:10:40.883Z)

I'm making a tiny readability cleanup in the new helper before final checks: wrapping a long docstring/temporary-path construction without changing behavior.

65. assistant (2026-07-02T07:10:52.006Z)

The code cleanup is done; I'm rerunning the focused test and style check against the touched files, then I'll do the final full coverage run before committing.

66. assistant (2026-07-02T07:10:58.297Z)

Focused verification is green. I'm running the full covered suite now so this branch has a fresh coverage number after the final source changes.

67. assistant (2026-07-02T07:13:18.536Z)

Full suite is green at 85.14% coverage against the 70% floor, so no coverage regression on the final P3 code. I'm going to re-verify branch and HEAD right before staging, per the repo's concurrency rules.

68. assistant (2026-07-02T07:13:22.708Z)

Branch and HEAD are cleanly based on `origin/master`, with no commits on the branch yet. I'm staging only the two P3 code/test paths now and committing them.

69. assistant (2026-07-02T07:13:31.025Z)

The P3 implementation commit is in place. I'm checking the post-commit branch state, then I'll

push and open the PR so I can add the exact PR number to the tracker in a follow-up amend.

70. assistant (2026-07-02T07:13:38.315Z)

The branch contains exactly one commit over `origin/master`. I'm checking whether the GitHub connector can create the PR cleanly here; if it can't, I'll use the already-authenticated `gh` fallback as the skill allows.

71. assistant (2026-07-02T07:13:41.376Z)

The connector is available this time. I'm pushing the branch first, then I'll try the connector-created PR with a ready-for-review body because you asked me to keep merging these.

72. assistant (2026-07-02T07:13:49.148Z)

The branch is on `origin/fix/source-cache`. I'm opening the PR now and will then update the live tracker with its number before the final push.

73. assistant (2026-07-02T07:14:01.006Z)

The connector hit the same repository visibility 404 as before, so I'm switching to the authenticated `gh` fallback. The branch is already pushed, so this should be straightforward.

74. assistant (2026-07-02T07:14:14.691Z)

PR #425 is open. I'm updating the live remediation tracker now so this row carries the actual PR number before CI/merge.

75. assistant (2026-07-02T07:14:19.771Z)

I'm marking only the P3 row as done in the tracker and leaving the later rows untouched.

76. assistant (2026-07-02T07:14:30.231Z)

The tracker edit is tiny, so I'm running the markdown link check before amending the commit.

77. assistant (2026-07-02T07:14:37.002Z)

The doc link check is clean. Before I amend, I'm doing the remaining repo-level lint/whitespace checks so the PR body can reflect the full verification set.

78. assistant (2026-07-02T07:14:44.089Z)

Ruff and whitespace checks are clean. I'm re-checking branch and HEAD immediately before staging the tracker file, then I'll amend the existing P3 commit.

79. assistant (2026-07-02T07:14:47.352Z)

The branch still has only the intended P3 commit. I'm staging just the tracker path and amending.

80. assistant (2026-07-02T07:15:01.400Z)

The amended commit now includes the tracker row. I'm checking the branch delta, then force-with-lease pushing the amended single-commit branch.

81. assistant (2026-07-02T07:15:05.813Z)

The branch contains one commit over `origin/master`; the ahead/behind is only from the intentional amend after PR creation. I'm using `--force-with-lease` so it will refuse if anything else touched the remote branch.

82. assistant (2026-07-02T07:15:21.072Z)

The amended branch is pushed. I'm updating PR #425's body with the final tracker/link-check verification, then I'll watch CI to merge only after it is green.

83. assistant (2026-07-02T07:15:38.898Z)

PR #425 is up to date with the final verification notes. I'm watching the GitHub checks now; if they pass, I'll squash-merge it and move to P4.

84. assistant (2026-07-02T07:18:37.247Z)

All GitHub checks on #425 are green, including full suite, coverage-backed CI, ruff, mypy baseline, markdown links, wheel smoke, and golden fixtures. I'm doing a quick PR metadata read, then I'll squash-merge and delete the branch.

85. assistant (2026-07-02T07:18:40.820Z)

The PR is clean, not draft, and the branch has exactly the one intended commit over `origin/master`. Merging it now.

86. assistant (2026-07-02T07:18:59.615Z)

PR #425 is merged and local `origin/master` fast-forwarded to `6c767b5`. I'm checking the checkout state now, then I'll refresh the handoff/current-state before starting P4.

87. assistant (2026-07-02T07:19:09.256Z)

The engine checkout is back on clean `master` at the merged P3 commit. I'm refreshing the external handoff files with the shipped state, then I'll branch from `origin/master` for P4.

88. assistant (2026-07-02T07:19:14.714Z)

I found the handoff still pointed at P3 as next. I'm updating the top live checkpoint to show #425 merged at `6c767b5`, with P4 as the next branch.

89. assistant (2026-07-02T07:19:36.946Z)

The handoff now points to P4. I'm also prepending the matching short checkpoint in `current-state.md` so the quick-start file stays honest.

90. assistant (2026-07-02T07:19:56.888Z)

Checkpoint files are refreshed. I'm moving into P4 now: fresh pull on engine master, explicit fetch, then a new branch from `origin/master` before reading the ffmpeg paths.

91. assistant (2026-07-02T07:20:14.471Z)

P4 branch is cut from `origin/master`. I'm searching the serving and media utility paths now to find subprocess calls that can run on request/proxy-serving code paths.

92. assistant (2026-07-02T07:20:23.634Z)

The audit points at two groups: job-path media subprocesses and the annotation proxy warmer. I'm checking the config model and current tests next so the timeout defaults fit existing patterns rather than becoming a scattered magic-number patch.

93. assistant (2026-07-02T07:20:41.098Z)

The live calls span modules where a typed app config is not always threaded through, so I'm checking call signatures before choosing the smallest coherent timeout API. The likely shape is a central ffmpeg helper with default/env override, plus config wiring where the compiler/proxy already has access.

94. assistant (2026-07-02T07:21:14.909Z)

The safest small fix is central timeout defaults in `backend.utils.ffmpeg`, then wiring each subprocess call to them and giving timeout-specific messages at the API-facing boundaries. I'm checking the job-worker comment too so we don't leave stale `ffmpeg` unbounded? prose behind.

95. assistant (2026-07-02T07:21:29.996Z)

I'm reading the existing ffmpeg resolver/video/compiler tests now. I want the new tests to assert the timeout contract without needing real ffmpeg, so coverage stays cheap and stable.

96. assistant (2026-07-02T07:21:41.643Z)

I'm adding the central timeout resolver first: defaults are bounded (`ffprobe` 60s, `ffmpeg` 1800s), with `RALLY\_FFPROBE\_TIMEOUT\_SEC` / `RALLY\_FFMPEG\_TIMEOUT\_SEC` overrides and `0` meaning `wait forever` for emergency compatibility.

97. assistant (2026-07-02T07:22:00.702Z)

Now I'm wiring the helper into the utility-level duration/chunk calls, with timeout-specific log messages and the same existing false/0.0 fallback behavior.

98. assistant (2026-07-02T07:22:14.100Z)

Next is the annotation proxy warmer, including the secondary P4 issue where a hung proxy transcode could keep a video stuck in the in-flight set until restart.

99. assistant (2026-07-02T07:22:32.030Z)

I'm wiring the compiler next. Per-clip trim and final concat timeouts become clean skip/failure messages, which flow through the existing `/api/compile` failed-result path instead of wedging the single worker.

100. assistant (2026-07-02T07:22:50.417Z)

I'm adding the same bounded `ffprobe` behavior to the validation path. A timeout there now becomes a normal validation failure rather than an unbounded subprocess.

101. assistant (2026-07-02T07:22:59.600Z)

Format validation is bounded; I'm doing PTS continuity and the resolution/fps fallback probe now.

102. assistant (2026-07-02T07:23:08.663Z)

PTS is bounded as a clean validation failure on timeout. I'm finishing the resolution/fps fallback probe.

103. assistant (2026-07-02T07:23:16.605Z)

I'm bounding the remaining process-path helpers now: Gemini debug trimming and audio extraction. These aren't the main compile wedge, but they were in the audit's serving-path list.

104. assistant (2026-07-02T07:23:28.544Z)

Gemini's only direct subprocess is the optional debug trim, so I'm adding the bounded timeout there and keeping its existing `fall back to full video` behavior on failure.

105. assistant (2026-07-02T07:23:36.965Z)

Audio extraction is next; if ffmpeg hangs, callers now get a concise RuntimeError instead of an unbounded wait.

106. assistant (2026-07-02T07:23:47.964Z)

The media calls are now bounded. I'm updating the job-worker risk note so it reflects that ffmpeg/ffprobe are covered by this P4 slice, while the broader per-job wall-clock timeout remains future work.

107. assistant (2026-07-02T07:24:02.322Z)

I'm adding focused tests around the timeout contract: resolver defaults/overrides, utility fallback behavior, compiler timeout errors, and proxy timeout fallback. These keep the change covered without invoking real ffmpeg.

108. assistant (2026-07-02T07:24:13.317Z)

The existing proxy tests mock subprocess with strict signatures, so I'm updating those stubs to accept `timeout=` and adding timeout-specific assertions.

109. assistant (2026-07-02T07:24:25.314Z)

I'm covering the video utility timeout fallbacks next.

110. assistant (2026-07-02T07:24:36.083Z)

I'm updating the annotation-proxy test doubles and adding a direct transcode-timeout fallback test.

111. assistant (2026-07-02T07:24:55.558Z)

I'm adding a compiler regression test for concat timeout mapping, which exercises the API-facing failure message the job worker will persist.

112. assistant (2026-07-02T07:25:17.994Z)

I'm adding a small timeout-specific test file for validator/audio boundaries, so the ?clean failure instead of hang? behavior is pinned where users would feel it.

113. assistant (2026-07-02T07:25:33.390Z)

The tests are in place. I'm scanning the touched serving-path modules for any remaining subprocess calls without `timeout=` before running them.

114. assistant (2026-07-02T07:25:39.598Z)

The scan shows only the expected touched calls; I'm reviewing the diff once before running tests, mostly to catch import formatting and message consistency.

115. assistant (2026-07-02T07:25:49.442Z)

The diff looks structurally right. I'm cleaning up one import style issue before running ruff so the formatter doesn't have to complain at me first.

116. assistant (2026-07-02T07:26:02.280Z)

I'm running the targeted timeout and media tests now, plus ruff over the touched files. If anything shakes loose, it should be in this batch.

117. assistant (2026-07-02T07:26:09.742Z)

The targeted timeout suite is green: 35 tests passed and ruff is clean for the touched paths. I'm broadening to the API/compile paths now because those have older subprocess mocks that can expose signature drift.

118. assistant (2026-07-02T07:26:39.561Z)

The broader API/media tests are green too: 80 passed, 1 skipped. I'm doing the full covered run

next so we have a fresh coverage number for P4, then I'll lint the whole repo and prepare the PR.

119. assistant (2026-07-02T07:28:58.929Z)

Full suite passed at 85.15% coverage, still well above the 70% floor. I'm running repo-wide lint, whitespace, and the local mypy check now; the mypy result should still hit the known `gcs.py` blocker.

120. assistant (2026-07-02T07:29:47.692Z)

Ruff and whitespace checks are clean; local mypy is still the known R1-27 `backend/storage/gcs.py:31` failure only. I'm checking the P4 diff and branch state now before staging.

121. assistant (2026-07-02T07:29:53.033Z)

Branch and HEAD are exactly right: P4 branch at `6c767b5`, no commits over `origin/master`, and the worktree only has the timeout files I touched. I'm staging explicit paths and committing the code/test slice.

122. assistant (2026-07-02T07:30:05.320Z)

The P4 code/test commit is in place. I'm checking the one-commit branch delta, then I'll push and open PR #426 (or whatever GitHub assigns) with the full coverage verification.

123. assistant (2026-07-02T07:30:09.533Z)

The branch has exactly one intended commit over `origin/master`. I'm pushing it now.

124. assistant (2026-07-02T07:30:22.656Z)

The branch is remote. I'm opening the PR via `gh` with the local verification, including the coverage number and the known local mypy caveat.

125. assistant (2026-07-02T07:30:39.542Z)

PR #426 is open. I'm updating the live tracker row for P4 with that PR number, then I'll amend and push the single-commit branch.

126. assistant (2026-07-02T07:30:48.704Z)

The tracker row is updated. I'm running the doc link check, then I'll re-verify branch/HEAD before amending.

127. assistant (2026-07-02T07:30:59.259Z)

Docs links are clean. I'm doing the required branch/HEAD re-check before staging the tracker path for the amend.

128. assistant (2026-07-02T07:31:02.743Z)

The branch still has just the intended commit, and the only unstaged change is the tracker row. I'm staging that explicit path and amending.

129. assistant (2026-07-02T07:31:11.172Z)

The amended P4 commit includes the tracker row. I'm confirming the single-commit branch then pushing with lease.

130. assistant (2026-07-02T07:31:15.446Z)

The branch is one amended commit over `origin/master`; the ahead/behind is only the intended

amend. I'm pushing with `--force-with-lease` now.

131. assistant (2026-07-02T07:31:22.432Z)

The amended branch is pushed. I'm updating PR #426's body to include the markdown link check before watching CI.

132. assistant (2026-07-02T07:31:36.783Z)

PR #426 is updated. I'm watching the GitHub checks and will only merge once the suite is green.

133. assistant (2026-07-02T07:34:45.994Z)

All GitHub checks on #426 are green, including full suite, ruff, mypy baseline, markdown links, wheel smoke, license scan, and golden fixtures. I'm doing the pre-merge metadata/delta check and then I'll squash-merge.

134. assistant (2026-07-02T07:34:50.196Z)

The PR is clean and the branch has exactly the one intended commit over `origin/master`. Merging #426 now.

135. assistant (2026-07-02T07:35:04.031Z)

PR #426 is merged and `origin/master` fast-forwarded to `c6bb4bf`. I'm confirming the checkout state and then refreshing the handoff/current-state before starting P5.

136. assistant (2026-07-02T07:35:08.719Z)

The engine checkout is clean on `master` at the merged P4 commit. I'm updating the live memory files with #426 and P5 as the next row.

137. assistant (2026-07-02T07:35:33.469Z)

Handoff is updated. I'm adding the matching short checkpoint to `current-state.md`.

138. assistant (2026-07-02T07:35:52.473Z)

Memory is current for P4. I'm moving on to P5: fresh pull, explicit fetch, and a new branch from `origin/master` before reading the scratch/retention paths.

139. assistant (2026-07-02T07:36:11.266Z)

P5 branch is cut from `origin/master`. I'm tracing where remote inputs create scratch directories and where cleanup tooling currently looks, so the fix covers the process/compile/assets/CLI/cloud marker cases without wandering into unrelated retention work.

140. assistant (2026-07-02T07:36:25.412Z)

The audit's P5 scope is clear: `acquire\_local\_input` makes `rally-input-\*` dirs; process, compile, assets, and CLI need scoped deletion, while `/api/source` was moved to `source\_cache` by P3; Cloud Run marker reads need `TemporaryDirectory`. I'm reading those call sites next.

141. assistant (2026-07-02T07:36:38.238Z)

I have the leak sites. I'm checking the rest of the asset/CLI control flow so cleanup happens after hashes/DB writes and never before consumers finish reading the scratch file.

142. assistant (2026-07-02T07:37:18.271Z)

The least disruptive shape is: keep `acquire\_local\_input()` returning a path, make it clean failed fetches itself, add a safe `cleanup\_acquired\_local\_input()` helper for callers? `finally` blocks, and use it in process/compile/assets/CLI. That avoids changing every caller into a

context manager and keeps local inputs as identity/no-op.

143. assistant (2026-07-02T07:37:30.288Z)

I'm checking the asset and compile tests in more detail before editing, because they already mock the storage seam and we can pin deletion there without introducing real cloud dependencies.

144. assistant (2026-07-02T07:37:39.488Z)

I'm adding the storage cleanup helper now. It only deletes directories with the `rally-input-` prefix and only for inputs that resolved remote, so local path behavior stays byte-for-byte.

145. assistant (2026-07-02T07:37:53.322Z)

I'm wiring that helper into `\_execute\_processing` first, after the report is emitted so observability still sees the local scratch path while it exists.

146. assistant (2026-07-02T07:38:02.013Z)

Process cleanup will run after the report hook, and is a no-op for local paths.

147. assistant (2026-07-02T07:38:08.976Z)

Next I'm putting compile's remote source cleanup in a `finally`, so success and clean failure both release the fetched video copy after stitching is done.

148. assistant (2026-07-02T07:38:19.877Z)

Cloud Run marker/report reads are simpler: I'm changing the one-off `cr-read-\*` directory to a `TemporaryDirectory` so it always disappears after JSON load.

149. assistant (2026-07-02T07:38:38.841Z)

For `/api/assets`, I'm scoping cleanup around the hash/store/register section. Uploaded local files are untouched; fetched `source\_uri` scratch copies are deleted even if a later store/register step raises.

150. assistant (2026-07-02T07:39:21.826Z)

The asset cleanup is in. I'm looking at the CLI branch tail so the remote CLI scratch copy is deleted after validation, segmentation, and report emission, including early returns.

151. assistant (2026-07-02T07:39:33.154Z)

I'm adding an outer `try/finally` around the CLI processing body. That way validation failures, indexing returns, and report emission all still release remote scratch.

152. assistant (2026-07-02T07:40:24.364Z)

Now I'm pinning the expected deletion behavior in the tests: process/compile/CLI scratch paths should be visible to consumers during work and gone after the job completes.

153. assistant (2026-07-02T07:40:32.536Z)

There is already an assets test file, good catch. I'm opening it and will add the remote-source scratch cleanup assertion there instead of overloading endpoint tests.

154. assistant (2026-07-02T07:40:37.171Z)

I'll start with storage unit tests: failed remote fetch cleanup, explicit successful cleanup, and no-op cleanup for local inputs.

155. assistant (2026-07-02T07:40:56.830Z)

Now I'm tightening the API/CLI tests so they assert scratch paths are deleted after consumers finish.

156. assistant (2026-07-02T07:41:14.986Z)

That patch missed one context chunk, so I'm splitting it into smaller surgical edits.

157. assistant (2026-07-02T07:41:18.722Z)

I've got the exact contexts now; applying the assertions in smaller chunks.

158. assistant (2026-07-02T07:41:24.907Z)

Compile/process post-job cleanup assertions are in. I'm adding the bucket scratch-dir list and CLI payload hash fix now.

159. assistant (2026-07-02T07:41:42.231Z)

For `/api/assets`, I'm adding the missing remote `source_uri` regression: fetch to hash, register URI as-is, then delete the scratch dir.

160. assistant (2026-07-02T07:41:54.918Z)

I'm pinning the Cloud Run marker read cleanup too: the temp dir should be gone after `_read_json` returns.

161. assistant (2026-07-02T07:42:15.621Z)

The P5 code and tests are in place. I'm running ruff on the touched files first; main.py indentation changes are the place I want a quick syntax/style check before tests.

162. assistant (2026-07-02T07:42:20.067Z)

Syntax and ruff are clean for the touched paths. I'm running the storage/API/compute tests that exercise the new cleanup behavior.

163. assistant (2026-07-02T07:42:48.602Z)

Targeted cleanup coverage is green: 107 passed, 1 skipped. I'm scanning for remaining `acquire_local_input` call sites so P5 doesn't miss another runtime path.

164. assistant (2026-07-02T07:42:56.784Z)

The remaining runtime `acquire_local_input` calls are exactly the four P5 targets. One subtle thing: reports are emitted while the scratch file exists, so I'm checking whether the report path is durable or could be written inside the soon-to-be-deleted scratch dir.

165. assistant (2026-07-02T07:44:06.703Z)

I'm picking up from the P5 branch that was already cut from `origin/master`. I'll first re-check the working tree and the handoff/current-state files so I don't accidentally work from stale assumptions, then I'll finish the scratch-retention fix and run the coverage gate before publishing.

166. assistant (2026-07-02T07:44:14.341Z)

The tree is on the intended P5 branch with only the scoped scratch-retention files modified. I'm now checking one subtle edge: whether deleting remote-input scratch after processing would also delete report artifacts that the API response points at.

167. assistant (2026-07-02T07:44:34.456Z)

Report emission does pass the resolved local scratch path to `obsreport`, and the comment says

reports default next to that video. I'm checking the local tests and config defaults before changing behavior, because this is exactly the kind of cleanup that can be correct for disk hygiene but wrong for response durability.

168. assistant (2026-07-02T07:44:51.742Z)

That edge is real: by default reports are next to the video, and for remote inputs the video is the temporary fetched copy. I'm going to keep the scratch deletion, but make remote-input reports land in a durable output-scoped report directory when no explicit `report.output\_dir` is configured.

169. assistant (2026-07-02T07:45:37.536Z)

I'll add this as a narrow helper in `backend/main.py` and pin it with API tests: remote inputs get a durable default report directory under `output/reports`, while explicit report configuration and local inputs keep existing behavior.

170. assistant (2026-07-02T07:45:55.650Z)

I'm editing now: adding the durable report-config helper, using it in API and CLI report emission, and adding one early-failure cleanup path so a fetched remote file is removed even if hashing raises.

171. assistant (2026-07-02T07:46:54.054Z)

The durability fix is in place now. I'm running the scoped lint/compile/tests again so any regression from that report-path correction shows up before I spend cycles on the full suite.

172. assistant (2026-07-02T07:46:57.491Z)

Scoped lint and syntax are clean. Next I'm running the targeted test group that covers the four scratch paths in this PR: process, compile, asset ingest, and cloud marker reads.

173. assistant (2026-07-02T07:47:30.951Z)

Targeted coverage is solid: `109 passed, 1 skipped`. I'm doing a quick diff read now, then I'll run the full covered suite to confirm the repo-level percentage doesn't regress.

174. assistant (2026-07-02T07:47:58.918Z)

One more small cleanup-hardening tweak: after a remote fetch succeeds, a database lookup failure would currently skip the main `finally`. I'm covering that narrow gap too so the scratch copy has a clear owner from fetch through job completion.

175. assistant (2026-07-02T07:48:05.147Z)

I'm also making the report helper read storage config correctly from either typed config or the legacy dict shape; it's a tiny compatibility polish while the function is fresh.

176. assistant (2026-07-02T07:48:13.511Z)

The compatibility polish is in; I'm rerunning the scoped gate after those last small edits.

177. assistant (2026-07-02T07:48:17.259Z)

Lint and syntax are still clean. I'll rerun the targeted tests once more, then full coverage.

178. assistant (2026-07-02T07:48:51.693Z)

Scoped tests passed again. Now I'm running the full coverage suite with the repo floor (`--cov-fail-under=70`) so we have the actual non-regression number for the PR.

179. assistant (2026-07-02T07:51:26.723Z)

Full suite is green, but the coverage percentage came back `85.13%`, which is above the 70 floor but a hair below the last checkpoint's `85.15%`. I'm going to add a tiny regression test for the new early-failure cleanup branch so the PR doesn't merely pass the floor; it should hold the line.

180. assistant (2026-07-02T07:51:37.128Z)

I found a clean coverage target: the branch that deletes a fetched remote file when hashing itself fails. I'm adding that regression test now; it also strengthens the P5 behavior rather than padding coverage for its own sake.

181. assistant (2026-07-02T07:51:48.588Z)

The extra regression test is in. I'm running the scoped gate again, then I'll rerun the full covered suite and compare the percentage.

182. assistant (2026-07-02T07:52:18.850Z)

The new test hit the submit-time segmenter validation before it reached hashing. I'm wiring the same fake segmenter used by the nearby process tests so the failure happens exactly where intended.

183. assistant (2026-07-02T07:52:27.743Z)

That fix is in; I'm rerunning the scoped tests now.

184. assistant (2026-07-02T07:52:56.744Z)

Scoped tests are back to green (`110 passed, 1 skipped`). I'm rerunning the full covered suite now; this is the one that determines whether coverage held at the repo level.

185. assistant (2026-07-02T07:55:22.770Z)

Full covered suite is back at `85.15%`, matching the last checkpoint and above the 70 floor. I'm moving through the remaining repo gates now: full ruff, mypy (expected known `gcs.py` failure until the next row), and diff whitespace.

186. assistant (2026-07-02T07:56:03.014Z)

The remaining gates look as expected: full `ruff` is clean; `git diff --check` is clean; local mypy hits only the known next-row `backend/storage/gcs.py:31` error. I'm checking the live tracker now so P5's row can be updated cleanly once the branch is published.

187. assistant (2026-07-02T07:56:13.138Z)

The tracker still has P5 as Todo, which is right until the PR exists. I'm doing the final pre-commit whitespace/status checks now, then I'll commit the code branch with explicit staging.

188. assistant (2026-07-02T07:56:21.092Z)

Pre-commit check is clean: branch is `fix/remote-scratch-retention`, HEAD is the P4 merge base, and there are no commits ahead yet. I'm staging only the seven P5 files explicitly and committing.

189. assistant (2026-07-02T07:56:43.059Z)

The code commit is made as `d21bc81`. I'm checking the branch delta now to make sure only that commit is ahead before I publish.

190. assistant (2026-07-02T07:56:51.577Z)

Branch delta is exactly one commit ahead of `origin/master`. I'm pushing it and opening the P5 PR with the verification results, including the coverage percentage and the known mypy blocker.

191. assistant (2026-07-02T07:57:46.728Z)

PR #427 is open. I'm adding the tracker update to the same branch now, then I'll run the docs link check and amend the commit so the PR carries its own plan status.

192. assistant (2026-07-02T07:57:55.400Z)

The tracker now points at #427. I'm validating docs links and whitespace before amending.

193. assistant (2026-07-02T07:58:01.355Z)

Docs validation is clean. I'm re-verifying branch and HEAD, then I'll stage only the tracker doc and amend the PR commit.

194. assistant (2026-07-02T07:58:09.489Z)

The amended commit is `bfa8d2b`. I'm force-pushing with lease and then I'll update the PR body to include the docs link check.

195. assistant (2026-07-02T07:58:30.840Z)

PR #427 now includes the tracker update and full verification notes. I'm checking CI; if it goes green, I'll merge it and then move on to the next row.

196. assistant (2026-07-02T07:58:36.221Z)

GitHub hasn't reported checks yet for the just-updated branch. I'll give Actions a moment to attach, then watch the check rollup rather than guessing.

197. assistant (2026-07-02T07:59:03.792Z)

CI is attached and running; early lint/docs/leak checks are already green. I'm watching the remaining jobs now.

198. assistant (2026-07-02T08:01:30.047Z)

All #427 GitHub checks are green, including the full suite. I'm merging it now and deleting the remote branch, then I'll sync master locally before starting the real `gcs.py` mypy fix.

199. assistant (2026-07-02T08:01:46.973Z)

#427 merged as `19ebe5a`. I'm checking local branch state now; `gh` fast-forwarded the checkout during merge, so I want a clean handoff to `master` before the next branch.

200. assistant (2026-07-02T08:02:03.110Z)

P5 is merged and local `master` is clean at `19ebe5a`. I'm refreshing the handoff/current-state memory now so the next session starts from the actual shipped state, then I'll cut the `gcs.py` mypy fix branch fresh from `origin/master`.

201. assistant (2026-07-02T08:02:30.154Z)

The handoff now reflects #427 as merged. I'm prepending the shorter current-state checkpoint next.

202. assistant (2026-07-02T08:02:40.528Z)

That patch didn't prepend the new checkpoint cleanly; I'm correcting it so the current-state top points at P5 merged, not the previous P4 pickup.

203. user (2026-07-02T08:02:58.011Z)

can you also make sure the project tracker is also updated

204. assistant (2026-07-02T08:03:15.727Z)

Yep, I updated the live tracker in the merged #427 commit: P5 is now `Done | #427` in `docs/KHELSUTRA\_GURU\_HEALTH\_REMEDIATION\_PLAN.md`. I'm verifying that directly on `master` before I branch for the next item, so the next PR starts from the real updated tracker.

205. assistant (2026-07-02T08:03:23.541Z)

Confirmed: `master` is up to date and the tracker row says `P5 ... Done | #427`. I'm moving to the next small PR now: the real `backend/storage/gcs.py` mypy fix, cut fresh from `origin/master`.

206. assistant (2026-07-02T08:03:33.290Z)

The mypy-fix branch is now fresh off `origin/master`. I'm reading the GCS backend and its tests next, then I'll make the smallest type/runtime-safe import change.

207. assistant (2026-07-02T08:03:52.636Z)

The failure is exactly the misplaced ignore: mypy reports the `from google.cloud import ...` line, while the ignore sits on the imported name. I'm switching that lazy import to `import\_module("google.cloud.storage")` and adding a fake-module test so the no-injected-client path is covered without requiring GCP credentials.

208. assistant (2026-07-02T08:03:56.784Z)

I'm making the import change now, plus the offline fake-client test for the lazy real-client path.

209. assistant (2026-07-02T08:04:12.966Z)

The GCS import is now dynamic and covered by an offline test. I'm running the focused storage test, then full mypy to confirm the known blocker is really gone.

210. assistant (2026-07-02T08:04:53.628Z)

Good, the R1-27 failure is gone: full `python -m mypy backend training --no-incremental` now succeeds locally with google-cloud-storage installed. I'm running full coverage next to guard the no-regression requirement.

211. user (2026-07-02T08:07:17.679Z)

can you pause after the current fix is merged

212. assistant (2026-07-02T08:07:30.598Z)

Got it. I'll finish and merge this `gcs.py` mypy fix, update the tracker for that PR, refresh memory, and then stop there rather than starting the next row.

213. assistant (2026-07-02T08:07:40.107Z)

The remaining local gates are clean too: full `ruff`, full coverage at `85.16%`, and full mypy now passes. I'm committing just the GCS import fix and its storage test next.

214. assistant (2026-07-02T08:07:51.512Z)

Committed as `a4f0524`. I'm verifying the branch delta is only this commit, then I'll push/open the PR and update the tracker row with that PR number.

215. assistant (2026-07-02T08:08:04.268Z)

Branch is pushed. Opening the PR now with the local verification, especially the full mypy pass and coverage number.

216. assistant (2026-07-02T08:08:15.759Z)

PR #428 is open. I'm updating the project tracker row for the `gcs.py` mypy fix to `Done | #428`, then I'll amend and push that into the same PR.

217. assistant (2026-07-02T08:08:25.166Z)

Tracker row P1 now points at #428 locally. I'm running the docs link check and whitespace check, then amending the PR commit.

218. assistant (2026-07-02T08:08:31.501Z)

Docs are clean too. Re-verifying branch and HEAD before staging the tracker doc, then I'll amend.

219. assistant (2026-07-02T08:08:43.242Z)

The amended commit is pushed. I'm updating the PR body with the tracker/docs verification and then watching CI.

220. assistant (2026-07-02T08:08:57.914Z)

#428 is ready and has the tracker update included. I'm checking CI status now; once it's green I'll merge and pause as requested.

221. assistant (2026-07-02T08:09:03.396Z)

Checks haven't attached yet after the force-push. I'll wait briefly for Actions to pick up the branch, then watch the full rollout.

222. assistant (2026-07-02T08:09:29.690Z)

CI is underway and the early jobs are green. I'm watching the rest through to completion.

223. assistant (2026-07-02T08:09:56.326Z)

GitHub's GraphQL call hiccuped mid-watch, not a CI failure. I'm retrying the check watch.

224. assistant (2026-07-02T08:11:58.933Z)

All #428 checks are green, including CI mypy and the full offline mocked suite. I'm merging it now, then I'll update memory and stop.

225. assistant (2026-07-02T08:12:18.815Z)

#428 merged as `d00b4db`; local checkout is fast-forwarded to `master`. I'm doing the final status check and refreshing memory with P1/R1-27 complete, then I'll pause.

226. assistant (2026-07-02T08:12:25.771Z)

Final status is clean: `master` is at `d00b4db`, and the tracker rows show `P1 Done #428` and `P5 Done #427`. I'm just refreshing the handoff/current-state now so the pause point is honest.

227. assistant (2026-07-02T08:13:21.952Z)

Paused here as requested.

Merged:

- #427 P5 scratch/retention as `19ebe5a`
- #428 R1-27/P1 `gcs.py` mypy fix as `d00b4db`

Tracker is updated:

- P5 is `Done | #427`

- P1 is `Done | #428`

Verification for #428:

- Full mypy now passes locally: `Success: no issues found in 141 source files`
- Full covered suite: `1570 passed, 11 skipped`
- Coverage: `85.16%` vs floor `70`, no regression
- GitHub CI all green

I also refreshed `session-handoff.md` and `current-state.md`. I did not start P6.

228. user (2026-07-02T08:20:21.463Z)

can you show me a table of the pending items left from the audit?

229. assistant (2026-07-02T08:20:30.111Z)

I'll pull the engine repo first per the freshness rule, then read the live remediation tracker and summarize only the rows that aren't done.

230. assistant (2026-07-02T08:20:38.163Z)

Repo is fresh and clean on `master`. I'm reading the tracker rows now and will collapse them into a pending-only table.

231. assistant (2026-07-02T08:21:00.243Z)

As of fresh `master` (`d00b4db`), these are the remaining non-done tracker rows from [KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN.md](C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/docs/KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN.md:143):

ID	Wave	Repo	Status	Pending item	Gate / note
P6	2	engine	Todo	Stream `GDriveApiStorage.fetch_to_local`; safe overwrite/update semantics	No gate
P7	2	engine	Todo	Cloud Run cancellation or late-marker salvage on poll timeout	Real cloud verification owner-gated
P8	2	engine / khelsutra docs	Todo	Upload-retention scheduling docs/unit/timer wrapper	Live infra enablement owner-gated
P9	3	engine	Todo	Pin `trackers` to immutable SHA; dedupe install literals	No gate
P10	3	engine	Todo	Update `obsreport` URL to `Khelsutra/` and immutable ref	No gate
P11	3	sports-obsreport	Todo	Add CI and fix `_plural("match", 3)`	No gate
P12	3	sports-data-collector	Todo	Add ruff + coverage gate after fixing ruff findings	Do not clobber dirty `.gitignore`
P13	3	sports-data-collector	Todo	Resolve stale `.gitignore` artifact hunk deliberately	Minor owner direction
P14	3	khelsutra	Todo	Add Playwright CUJ CI + mock/contract drift guard	No gate
P15	4	engine + khelsutra	Todo	Refresh engine `backend/ui/` bundle; staleness check/doctor warning	P14 preferred
P16	4	khelsutra	Todo	Refresh-resumable SPA ingest jobs	Needs screencast evidence
P17	4	khelsutra	Todo	Keep uploaded filename visible	Screenshot/screencast if visible
P18	4	khelsutra	Todo	Hide/replace demo rally grid during real job	Needs screencast evidence
P19	4	khelsutra	Todo	Handle HTML/proxy 5xx bodies safely	No gate
P20	4	khelsutra	Todo	Web Share API file-share for compiled reels	Mobile-capable validation notes
P21	4	khelsutra	Todo	Replace stale SPA API line comments with anchors	No gate
P22	4	rally-annotator	Todo	Canonical remote URL + remaining reload/persistence E2E	No gate
P23	5	engine	Todo	Close/re-cut stale `#390` main.py split	After P1-P8 settle
P24	5	engine	Todo	Remove `backend.main` shim cycles; lift mypy quarantine	After P23
P25	5	engine	Todo	Split `Database` behind domain facade	After P23
P26	6	mixed	Todo	Update/close stale GitHub issues	Some owner disposition

| P27 | 6 | engine / non-git | Todo | Back up Roora pilot non-video artifacts; update
`DATA\_IN\_GCS.md` | Video upload owner-gated |
| P28 | 6 | sports-data-collector / non-git | Blocked | Hash six owned golden source videos from
`F:` | Needs `F:` mounted; upload approval |

232. user (2026-07-02T08:44:25.802Z)

There were some fixes in the other session. Can you resync to head and update the pending items list and go on fixing the items which are in the repo khelsutra, rally-annotator, sports-data-collector or any non git folder. Don't make changes to the badminton-highlight@indexer repo and make sure all others are fixed.

233. assistant (2026-07-02T08:44:42.272Z)

I'll resync every non-engine repo before reading, then I'll re-read the live tracker from engine read-only. I'll avoid editing `badminton-highlight-indexer` entirely and start fixing the non-engine rows one PR-sized slice at a time.

234. assistant (2026-07-02T08:45:12.836Z)

Engine advanced from `d00b4db` to `65983ac`; I'll treat it as read-only and use its tracker as the source of truth. The listed non-engine repos are fresh; I'm also checking `sports-obsreport` because the tracker may still have a row there, but I won't touch anything outside your requested scope without calling it out.

235. assistant (2026-07-02T08:45:24.581Z)

`badminton-highlight-indexer`, `khelsutra`, `rally-annotator`, and `sports-obsreport` are clean at head. `sports-data-collector` still has the pre-existing dirty `.gitignore` hunk, so I'll avoid staging or clobbering it unless we deliberately address P13.

236. assistant (2026-07-02T08:45:39.372Z)

After resync, engine-side P6 and P7 are already marked done by the other session. The remaining requested-scope rows are mainly `khelsutra` P14-P21, `sports-data-collector` P12/P13/P28, `rally-annotator` P22, and non-git P27/P28. I'll start with a non-UI, small scoped item: rally-annotator P22.

237. assistant (2026-07-02T08:45:47.069Z)

For P22 I'm cutting a fresh branch from `origin/main` in `rally-annotator`, then I'll look for stale repo URLs and existing reload/persistence coverage before editing.

238. assistant (2026-07-02T08:46:14.544Z)

P22 has two parts. I found the local `origin` still points at `avidullu/rally-annotator`, and the VLC metadata still has that old URL, while README badges are already canonical. The web E2E suite already records screencasts; I'm checking whether it actually exercises reload persistence or only storage writes.

239. assistant (2026-07-02T08:46:36.621Z)

I'm going to fix both concrete P22 surfaces: canonicalize the stale GitHub URL in the VLC extension metadata and add a Playwright reload check that proves saved rows survive page refresh and resume numbering. I'll also update this checkout's `origin` URL to the canonical org, since that part is local Git config rather than a commit.

240. assistant (2026-07-02T08:46:41.584Z)

Local `origin` now points to `Khelsutra/rally-annotator`. I'm applying the code/test patch for the committed part.

241. assistant (2026-07-02T08:47:18.590Z)

The P22 patch is in. I'm running the rally-annotator gates now: Lua syntax/dialog tests, web typecheck/unit coverage/build, then the Chromium E2E if the local browser dependencies are present.

242. assistant (2026-07-02T08:47:36.561Z)

Whitespace is clean. I'm checking whether the local machine has the Lua and Node toolchains ready; the web E2E will be run after the extension build.

243. assistant (2026-07-02T08:47:49.026Z)

Lua 5.1 isn't installed on this Windows side, so I can't run the VLC dialog suite locally unless WSL has it. Node is available but `node\_modules` isn't present, so I'm doing a clean `npm ci` before the web gates.

244. assistant (2026-07-02T08:48:15.507Z)

WSL does have Lua 5.1, so I'll run the VLC checks there. The web dependencies are installed; next up are typecheck, coverage, build, and then the real Chromium E2E screencast run.

245. assistant (2026-07-02T08:48:28.400Z)

All non-E2E gates passed: Lua dialog suite `88 passed`, web typecheck, coverage `97.85%`, and extension build. I'm running the real Chromium E2E now; it should also produce updated per-locale screencasts under `web/e2e-results/screencasts`.

246. assistant (2026-07-02T08:48:41.061Z)

The E2E failure is environmental: Playwright Chromium wasn't installed locally. I'm installing Chromium with the repo's own Playwright command and rerunning the E2E, so we get real screencast evidence instead of stopping at that setup gap.

247. assistant (2026-07-02T08:49:17.175Z)

Chromium is installed now. I'm rerunning the E2E and will inspect/cite the produced screencast paths if it passes.

248. assistant (2026-07-02T08:49:49.873Z)

E2E passed across all seven locales and produced the per-locale recordings. I'm checking the working tree and artifact paths now, then I'll commit only the source/test files.

249. assistant (2026-07-02T08:50:05.172Z)

The recordings are ignored as intended under `web/e2e-results/screencasts` (`en.webm`, `kn.webm`, `hi.webm`, `es.webm`, `da.webm`, `id.webm`, `te.webm`). No stale `avidullu/rally-annotator` references remain outside generated/ignored areas. I'm doing the pre-commit branch/HEAD check and committing just the two changed files.

250. assistant (2026-07-02T08:50:09.939Z)

Branch is correct and there are no extra commits ahead. I'm staging explicit paths only and committing.

251. assistant (2026-07-02T08:50:18.132Z)

Committed as `a33eda4`. I'll push and open the P22 PR with the verification plus screencast evidence paths.

252. assistant (2026-07-02T08:50:28.330Z)

Branch is pushed to the canonical org remote. Opening the PR now.

253. assistant (2026-07-02T08:50:47.422Z)

P22 PR is open as rally-annotator #28. I'm watching CI and will merge if it goes green. I'm not updating the engine tracker file because you explicitly said not to make changes in `badminton-highlight-indexer`.

254. assistant (2026-07-02T08:51:09.805Z)

The Lua CI check is green; the web extension CI is still running. I'll keep watching.

255. assistant (2026-07-02T08:52:08.872Z)

All rally-annotator #28 checks are green. I'm merging it and syncing local `main`.

256. assistant (2026-07-02T08:52:23.291Z)

P22 is merged in rally-annotator as `ef9e088`. I'm moving to the next requested-scope row, sports-data-collector P12. I'll keep the existing dirty `.gitignore` hunk untouched and stage explicit paths only.

257. assistant (2026-07-02T08:52:32.327Z)

First I'm refreshing sports-data-collector again, then branching from `origin/master` with the known dirty `.gitignore` left in place.

258. assistant (2026-07-02T08:52:37.564Z)

`master` is fresh; only the known `.gitignore` hunk is dirty. I'm cutting the P12 branch from `origin/master` now.

259. assistant (2026-07-02T08:52:51.087Z)

The branch is cut; the `.gitignore` WIP carried over but remains untracked by my work. I'm reading the collector's config, CI, and current lint/test state now.

260. assistant (2026-07-02T08:53:02.375Z)

Collector currently has pytest-only CI and no `ruff` dependency. I'm running current lint/coverage locally to see the exact fixes needed before wiring the gates.

261. assistant (2026-07-02T08:53:20.492Z)

Current collector coverage is `82%` with `289 passed, 2 skipped`; the missing part is lint gating. Ruff reports mostly unused imports plus a couple of small cleanup issues. I'm letting Ruff apply the safe fixes, then I'll manually review the remaining findings.

262. assistant (2026-07-02T08:53:34.064Z)

Ruff fixed the safe import/f-string issues and left three that need a human glance: two dead assignments in the crawler and one import-order issue in a test file. I'm reading those spots before editing.

263. assistant (2026-07-02T08:53:42.148Z)

The stat includes the pre-existing `.gitignore` WIP; I'm leaving that alone. I'm removing the two dead crawler assignments and moving the test import to the top so Ruff can pass without changing behavior.

264. assistant (2026-07-02T08:53:59.464Z)

The lint cleanups are in. Next I'm adding Ruff to the dev extra and changing CI from plain pytest to Ruff plus coverage with an 80% floor, which is below the current 82% but high enough to catch regressions.

265. assistant (2026-07-02T08:54:16.280Z)

The CI gate is wired. I'm running the exact new local gates now: `ruff check .` and the covered pytest command with the 80% floor.

266. assistant (2026-07-02T08:54:39.069Z)

Collector P12 gates pass locally: Ruff clean, coverage `81.94%` over an 80% floor with `289 passed, 2 skipped`. I'm checking the diff now and making sure the pre-existing `.gitignore` stays out of the commit.

267. assistant (2026-07-02T08:55:51.102Z)

I'll use the GitHub PR publishing workflow here because we're creating and merging small remediation PRs. I'm picking up from the resynced tracker state, keeping `badminton-highlight-indexer` read-only, and finishing the `sports-data-collector` gate PR next.

268. assistant (2026-07-02T08:56:10.488Z)

The collector branch has exactly the expected mixed tree: our lint/coverage edits plus the pre-existing `.gitignore` change. I'm going to refresh the remote ref, inspect the diff in small pieces, and then commit only the intended paths.

269. assistant (2026-07-02T08:56:26.535Z)

Remote `origin/master` still matches the branch base (`9ce5a20`), so there's no sibling commit riding along. The `.gitignore` hunk is still separate and I'll leave it untouched.

270. assistant (2026-07-02T08:56:35.188Z)

The ruff cleanup is small and mechanical: unused imports, two dead local assignments, and a CI/dev dependency update. The functional gate change is just adding `ruff` plus `pytest --cov=src --cov-fail-under=80` to CI.

271. assistant (2026-07-02T08:56:49.471Z)

I'm going to rerun the local gates and also compare coverage against a clean `origin/master` worktree. That gives us the `?floor is met?` and `?coverage didn't regress?` evidence instead of relying on vibes in a trench coat.

272. assistant (2026-07-02T08:57:09.257Z)

Current branch is green at 81.94% coverage. I'm spinning up a temporary detached worktree at `origin/master` now, purely for the baseline coverage comparison, then I'll remove it.

273. assistant (2026-07-02T08:57:38.315Z)

The baseline at `origin/master` is also 289 passed / 2 skipped, with exact coverage around 81.97%. Our branch is 81.94%, so I'm going to add one focused coverage bump before publishing; the rounded report is fine, but the exact denominator shifted enough that I don't want to hand-wave the regression.

274. assistant (2026-07-02T08:57:56.146Z)

I found several one-line gaps in already-tested core helpers. I'm going to cover the smallest useful one, rather than adding broad tests unrelated to this PR.

275. assistant (2026-07-02T08:58:14.094Z)

The coverage bump is a single parser edge case: invalid colon-delimited time strings. It's already in the nearby time parser test, and it covers the fallback path that was previously unexercised.

276. assistant (2026-07-02T08:58:18.877Z)

That test is in place. I'm rerunning lint and the coverage gate now, then I'll recheck the branch ancestry immediately before staging.

277. assistant (2026-07-02T08:58:39.607Z)

Coverage is now 82.02%, above the clean `origin/master` baseline, and the gate stays at 80%. I'm doing the final branch/HEAD verification and then staging only the intended files.

278. assistant (2026-07-02T08:58:44.939Z)

The branch is still exactly on `origin/master` with no local commits yet. I'm staging the workflow, pyproject, ruff cleanups, and the small parser test only; `.gitignore` stays out.

279. assistant (2026-07-02T08:58:57.257Z)

The staged set is clean and excludes the dirty `.gitignore` hunk. Committing this collector P12 fix now.

280. assistant (2026-07-02T08:59:09.040Z)

Commit `2d5436e` is ready. I'm checking the branch range one more time, then I'll push and open the PR with the verification notes included.

281. assistant (2026-07-02T08:59:14.190Z)

The branch range contains only our collector gate commit, and the only unstaged work remains the pre-existing `.gitignore`. Pushing now.

282. assistant (2026-07-02T08:59:22.220Z)

The branch is on GitHub. I'm opening the PR with the exact coverage comparison and the local checks so CI has the same contract documented.

283. assistant (2026-07-02T08:59:34.751Z)

Collector PR #58 is open. I'm watching CI now; if it goes green, I'll merge it and then move to the next allowed repo.

284. assistant (2026-07-02T09:00:04.114Z)

CI passed on #58. I'm merging it with squash and deleting the remote branch, then I'll fast-forward local `master` while preserving the untouched `.gitignore` WIP.

285. assistant (2026-07-02T09:00:19.197Z)

PR #58 is merged and the local checkout fast-forwarded to `01ceda9`. I'm checking the collector state, then I'll refresh the allowed-repo pending list and start the next fix.

286. assistant (2026-07-02T09:00:37.636Z)

Collector P12 is merged. I'm going back to the read-only engine tracker now to refresh the pending list, then I'll pull `kheksutra` and branch from its remote for the next allowed remediation row.

287. assistant (2026-07-02T09:00:46.409Z)

The shared memory is still stale relative to the PRs we just merged, so I'll treat the live tracker plus remote heads as truth. The `kheksutra` repo is freshly pulled; next I'm reading the plan rows for the remaining allowed repos.

288. assistant (2026-07-02T09:00:58.652Z)

The read-only tracker still says P12/P22 are Todo because it lives in the engine repo, which you told me not to change. Ground truth after the PRs: P12 and P22 are merged; the next allowed unblocked row is Khelsutra P14.

289. assistant (2026-07-02T09:01:02.791Z)

Khelsutra is clean and freshly pulled. For P14 I'm cutting a new branch directly from `origin/master`, then I'll inspect the existing web test/e2e scripts and CI shape before editing.

290. assistant (2026-07-02T09:01:22.281Z)

P14 already has the raw ingredients: a hermetic mock engine, Playwright CUJ scripts, and screencast output conventions. I'm checking the existing GitHub workflow and the CUJ coverage gaps now, then I'll wire the CI profile without disturbing the artifact convention.

291. assistant (2026-07-02T09:01:36.986Z)

The current CI builds and unit-tests the SPA but never runs the hermetic Playwright CUJ. I'm going to add a CI-specific Playwright mode, then add a lightweight contract guard so mock-engine response shapes don't drift silently from the SPA client expectations.

292. assistant (2026-07-02T09:02:59.752Z)

I'll keep the local screencast behavior intact. The CI path will build the SPA, run a Node mock-contract smoke against every API route the client relies on, then run the Playwright CUJs with CI-friendly video/trace retention on failure.

293. assistant (2026-07-02T09:03:07.408Z)

The package script is in. Next I'm adapting the Playwright config so normal local screencasts still record, while CI runs faster and preserves artifacts only when something fails.

294. assistant (2026-07-02T09:03:17.587Z)

Now I'm adding the contract smoke: it boots the same mock engine the CUJs use and probes the response shapes for the SPA client's routes before Playwright starts.

295. assistant (2026-07-02T09:03:56.282Z)

The contract smoke covers capabilities/config/videos/process/query/compile/assets/annotations/upload/source/reels. I'm wiring that into the web CI job now with Playwright browser install and failure artifacts.

296. assistant (2026-07-02T09:04:04.648Z)

CI is wired. I'm adding the short doc note so future agents know `e2e:ci` is the non-recording gate and `screencast` remains the local evidence path.

297. assistant (2026-07-02T09:04:15.741Z)

The edits are all in CI/e2e harness territory, but I'll still run the same web gates: clean install, typecheck, coverage, and the new CI e2e profile. If Chromium is missing locally, I'll install it and rerun.

298. assistant (2026-07-02T09:04:28.236Z)

Install is clean, no vulnerabilities. Running typecheck and coverage now, then the Playwright CI profile.

299. assistant (2026-07-02T09:04:56.796Z)

Typecheck and coverage both passed. I'm installing/refreshing Chromium for the local Playwright

run, then I'll run the new `e2e:ci` path exactly as CI will.

300. assistant (2026-07-02T09:05:58.118Z)

The local e2e run passed all four CUJs and produced raw screencast videos in the existing sibling artifact folder. I'm also running it once with `CI=true` and the CI `SCREENCAST\_DIR` override, because GitHub Actions will exercise that exact mode.

301. assistant (2026-07-02T09:06:48.891Z)

The real CI mode passes too. I'm listing the local screencast artifacts now, then I'll do diff hygiene and commit P14.

302. assistant (2026-07-02T09:07:03.256Z)

The CI-mode run created an empty/local `web/e2e-results/` artifact directory. I'm checking ignore rules; if it isn't already ignored, I'll add a narrow ignore so CI/debug artifacts don't pollute future working trees.

303. assistant (2026-07-02T09:07:08.083Z)

`web/e2e-results` is a CI/local artifact dir and wasn't ignored. I'm adding that narrow ignore in `web/.gitignore` so the new CI artifact path stays out of commits.

304. assistant (2026-07-02T09:07:26.872Z)

The artifact ignore is clean now, and the local screencast evidence is present: annotate, reopen/build/download, and store-medium `.webm` recordings in the sibling `khelsutra-screencast` folder. I'm reviewing the patch, then I'll stage explicit paths.

305. assistant (2026-07-02T09:07:49.131Z)

The P14 web gates are green. Since the repo's CI also has site/layout jobs, I'm going to run those local equivalents too before I publish the PR. It's a little extra time, but it keeps the merge stamp honest.

306. assistant (2026-07-02T09:07:55.511Z)

The first shell-script parse check used Windows absolute paths, which Git Bash mangled. I'm rerunning it with repo-relative POSIX paths, matching the Linux CI shape.

307. assistant (2026-07-02T09:09:15.537Z)

The full local CI shape is green now: site tests, shell parse checks, layout/render/a11y, plus the web gates. I'm doing the branch/HEAD verification and staging only the P14 files.

308. assistant (2026-07-02T09:09:19.044Z)

The branch is still exactly on `origin/master` with no extra commits. I'm staging six explicit files now; no generated artifacts are included.

309. assistant (2026-07-02T09:09:25.066Z)

Staged set looks right. Committing the P14 Khelsutra CI/CUJ gate now.

310. assistant (2026-07-02T09:09:33.523Z)

Commit `2c9cf06` is ready. I'm checking that the branch contains only this commit, then I'll push and open the PR.

311. assistant (2026-07-02T09:09:39.852Z)

The branch range is clean: one commit, target `avidullu/khelsutra` `master`. Pushing now.

312. assistant (2026-07-02T09:09:46.110Z)

Branch is pushed. I'm opening the PR with the coverage, full local CI, and screencast evidence paths recorded.

313. assistant (2026-07-02T09:10:04.859Z)

Khelsutra PR #143 is open. I'm watching CI now; if it passes, I'll merge it and then continue to the next allowed Khelsutra row.

314. assistant (2026-07-02T09:12:29.417Z)

All checks on #143 are green: site, web, layout, and the Cloudflare Workers build. I'm merging it now and fast-forwarding local `master`.

315. assistant (2026-07-02T09:12:43.723Z)

P14 is merged as `7b6544f`. I'm checking the checkout state, then I'll start the next Khelsutra wave row from fresh `origin/master`.

316. assistant (2026-07-02T09:12:51.434Z)

Next allowed row is P16. I'm skipping P15 for now because it requires modifying the engine's vendored `backend/ui/` bundle, and you explicitly told me not to change `badminton-highlight-indexer`. P16 is Khelsutra-only, so I'll branch it fresh.

317. assistant (2026-07-02T09:13:06.870Z)

Branch `fix/job-resume-progress` is fresh from merged `origin/master`. I'm reading the ingest flow and current job polling now so I can make the resume behavior fit the existing state model.

318. assistant (2026-07-02T09:13:17.545Z)

The current hook only keeps polling in memory; a refresh loses the job. I'm going to store an active job record (`job\_id`, URI, cid, start time, locale/compute), support `?job=...` resume on app load, and surface elapsed/heartbeat text in the existing ingest status chip.

319. assistant (2026-07-02T09:14:54.230Z)

I'm going to refactor `useIngestJob` first. The key behavior: once the engine returns a `job\_id`, the SPA writes it to `?job=` and `localStorage`; on refresh, the hook resumes polling that same job instead of making the user submit again.

320. assistant (2026-07-02T09:15:57.228Z)

The hook persistence is in place. I'm updating the panel copy now so the operator sees that the job resumed, how long it has been running, and that the page can be safely refreshed.

321. assistant (2026-07-02T09:16:07.190Z)

The UI now has a heartbeat line. I'm adding the matching catalog entries across all six locales so the i18n parity and placeholder guards stay green.

322. assistant (2026-07-02T09:16:57.538Z)

I noticed one edge while writing the resume test: an immediate resume poll should back off to the normal one-second cadence after the first check, not keep scheduling zero-delay polls. I'm tightening that now.

323. assistant (2026-07-02T09:17:02.833Z)

Now I'm adding unit coverage for the two important resume contracts: new submits write `?job=/localStorage`, and a refresh resumes a saved job without calling `/api/process` again.

324. assistant (2026-07-02T09:17:25.916Z)

The hook tests now pin persistence and resume. I'm adding a panel-level assertion that the heartbeat and reload-safe hint are actually visible to the user.

325. assistant (2026-07-02T09:17:32.727Z)

Next I'm adding the Playwright CUJ for the refresh behavior itself: submit, observe `?job=`, reload, resume, then land on the loaded video.

326. assistant (2026-07-02T09:17:58.572Z)

For the resume CUJ to be deterministic, I'm making the mock keep the special `resume\_me` job running for one extra poll. Normal CUJs and the contract smoke keep their existing two-poll behavior.

327. assistant (2026-07-02T09:18:09.437Z)

The mock now gives the resume scenario enough time to show the resumed heartbeat. I'm running the focused hook/panel/i18n tests before the full web suite.

328. assistant (2026-07-02T09:18:26.078Z)

The focused tests pass, but one test left `?job=` in jsdom and the next test auto-resumed noisily. I'm cleaning the test setup so URL state is reset between component tests too.

329. assistant (2026-07-02T09:18:36.321Z)

Test URL cleanup is fixed. I'm rerunning the focused set once more, then I'll move to typecheck/full coverage/e2e.

330. assistant (2026-07-02T09:18:47.781Z)

Focused tests are green. I'm running the full web typecheck and coverage gate now; this is the coverage comparison that matters for this SPA change.

331. assistant (2026-07-02T09:19:09.993Z)

Typecheck passed, but the full coverage gate caught a real branch-coverage drop: 86.45% vs the 87% floor. That's exactly why we run it. I'm going to add a couple of focused branch tests around resume fallbacks and poll-error retention instead of touching thresholds.

332. assistant (2026-07-02T09:20:44.066Z)

I'm picking up the in-progress `khealsutra` P16 work rather than touching the engine tracker. First I'll re-orient the local state carefully, then I'll finish the coverage gap and only stage the intended paths.

333. assistant (2026-07-02T09:21:12.683Z)

The allowed repos are all still at the expected merged heads so far. I'm doing the explicit fast-forward pulls now, including the dirty `sports-data-collector` checkout, and I'll keep its `.gitignore` hunk out of any future staging.

334. assistant (2026-07-02T09:21:26.521Z)

I'm also going to read the live audit tracker read-only so the pending table reflects the current document, but I won't edit anything inside `badminton-highlight-indexer`.

335. assistant (2026-07-02T09:21:38.960Z)

The engine docs moved under a folded/archive structure in the other session, so the old `live tracker` path is no longer present at head. I'll use the new audit report plus the folded plan

for a read-only pending table, and keep the actual fixes in the allowed repos.

336. assistant (2026-07-02T09:21:54.197Z)

The tracker at engine head now lives inside `docs/AUDIT\_REPORT\_khelsutra-guru\_2026-07-02.md`. It still doesn't know about our P12/P14/P22 merges, and because the latest instruction forbids engine edits, I'll treat that as a reporting-only delta for now. Next I'm tightening the P16 tests so the `khelsutra` branch coverage floor stays above 87%.

337. assistant (2026-07-02T09:22:00.776Z)

The current hook implementation is doing the right shape of work, but it introduced a few defensive branches around URL/localStorage failures and resume cases. I'm going to use the coverage table to target those branches instead of shaving the feature down.

338. assistant (2026-07-02T09:22:29.222Z)

Full coverage ran with thresholds neutralized; the useful part is the generated coverage artifacts. I'm reading the summary now so the next patch is just tests around the uncovered resume/error branches.

339. assistant (2026-07-02T09:22:36.365Z)

The first diagnostic reporter didn't persist artifacts, so I'm switching to a JSON summary reporter. That should give us exact file-level branch numbers without touching source files.

340. assistant (2026-07-02T09:22:56.239Z)

The JSON report is available now. I'll inspect just the files with weaker branch coverage, then patch the tests and remove the temporary coverage directory before committing.

341. assistant (2026-07-02T09:23:17.937Z)

I've got the coverage culprit: `useIngestJob.ts` branch coverage is the main drag, not the UI wrapper. I'm adding focused hook tests for URL-only resume, invalid saved state, poll-error retention, and timeout retention, plus trimming one unreachable guard in the resume branch.

342. assistant (2026-07-02T09:23:23.437Z)

That unreachable guard is gone. Now I'm adding the tests that exercise the resume/error contract: no duplicate POSTs, saved jobs survive transient engine trouble, and stale polling times out without forgetting the job id.

343. assistant (2026-07-02T09:23:51.127Z)

The extra tests are in. I'm running the hook-focused set first so any fake-timer or jsdom weirdness shows up before the full gate.

344. assistant (2026-07-02T09:24:03.401Z)

The focused tests are green now. Because this PR changed root locale catalogs, I'm syncing the site's mirrored locale files before the full web/site gates, so the site drift test doesn't catch us later.

345. assistant (2026-07-02T09:24:06.962Z)

The site has a `copy:locales` script, so I'm using that rather than hand-editing the mirror files.

346. assistant (2026-07-02T09:24:15.739Z)

Locale mirrors are synced. I'm running the web CI-equivalent gate now: typecheck first, then coverage with the repo floor intact.

347. assistant (2026-07-02T09:24:39.915Z)

The web typecheck and coverage gate pass with the real thresholds back on. I'm doing the remaining web build and then the Playwright CUJ runs, including the new refresh-resume screencast.

348. assistant (2026-07-02T09:24:53.526Z)

Build is clean; Vite still prints the existing large-chunk warning. I'm checking the E2E scripts and then running the non-CI mode first so Playwright records the screencasts.

349. assistant (2026-07-02T09:25:00.794Z)

The harness records video outside CI, and the new test is in the same `\*.pw.mjs` set. I'll install/check Chromium and run `npm run e2e` for evidence, then the exact `e2e:ci` command with `CI=true`.

350. assistant (2026-07-02T09:25:06.875Z)

Chromium is present. Starting the recording E2E pass now; this one should drop fresh raw videos into `khelsutra-screencast/\_raw`.

351. assistant (2026-07-02T09:26:02.399Z)

The recording E2E pass is green: 5/5. I'm running the same gate in CI mode now, with failure artifacts contained under `web/e2e-results`.

352. assistant (2026-07-02T09:26:51.256Z)

CI-mode E2E is green too. I'm moving to the site gates now because the mirrored locale files changed: unit tests plus browser layout/render/all.

353. assistant (2026-07-02T09:28:02.674Z)

All site gates are green, including the locale mirror check. I'm cleaning the temporary diagnostic coverage folder now, then I'll do whitespace/status checks and collect the screencast paths for the PR.

354. assistant (2026-07-02T09:28:09.831Z)

Temporary diagnostics are cleared. I'm generating one quick JSON coverage summary so the PR can cite exact coverage numbers, then I'll remove that report too.

355. assistant (2026-07-02T09:28:29.537Z)

The JSON summary run is complete. I'm reading the exact totals and deleting the generated report so no diagnostic artifact can sneak into the PR.

356. assistant (2026-07-02T09:28:44.903Z)

Exact web coverage is now `94.08%` statements/lines, `87.42%` branches, `88.40%` functions, above the configured floors. I've removed the temp report and I'm running diff whitespace/status checks next.

357. assistant (2026-07-02T09:28:50.349Z)

Whitespace is clean and the only ignored artifacts are the expected local dependencies/build/E2E outputs. I'm expanding the screencast listing so the PR body has full evidence paths.

358. assistant (2026-07-02T09:28:59.710Z)

The refresh-resume screencast is present at `khelsutra-screencast/\_raw/refresh-resume.../video.webm`. I'm doing a code diff review now, then I'll re-verify branch/base before staging.

359. assistant (2026-07-02T09:29:11.563Z)

The diff review found one real edge: canceling while the initial ``POST /api/process`` is still pending could still start polling after the response returns. I'm fixing that by restoring a submit-run cancellation guard and adding a regression test.

360. assistant (2026-07-02T09:29:18.612Z)

The submit cancellation guard is back in place. I'm adding the test that holds the process response open, cancels, and proves no ``?job`` or saved job appears after the late response.

361. assistant (2026-07-02T09:29:33.196Z)

I'm tightening that test typing a bit before running it, so TypeScript doesn't need to pretend the mock response is a real browser ``Response``.

362. assistant (2026-07-02T09:29:41.154Z)

The cancellation regression test is in. I'm rerunning the focused hook/UI/i18n set, then I'll repeat the full coverage gate because the hook changed after the previous pass.

363. assistant (2026-07-02T09:29:53.402Z)

The regression test passes. I'm rerunning the full web gate now from scratch after that fix, then I'll rerun CI-mode E2E so the last verification matches the final code.

364. assistant (2026-07-02T09:30:16.181Z)

Full web typecheck, coverage, and build are green after the cancellation fix. I'm rerunning CI-mode Playwright now as the final browser gate.

365. assistant (2026-07-02T09:31:06.649Z)

CI-mode Playwright is green on the final code. I'm refreshing the exact coverage totals once more after the cancellation test so the PR body doesn't cite stale numbers.

366. assistant (2026-07-02T09:31:27.914Z)

The exact summary is generated. I'm reading it and then deleting it; after that I'll do the final diff check and stage.

367. assistant (2026-07-02T09:32:52.353Z)

I'm picking back up at the in-progress ``khelsutra`` P16 branch, then I'll publish and merge that one PR and pause as requested. I'll keep the engine repo read-only and report the tracker delta instead of editing its files.

368. assistant (2026-07-02T09:32:56.951Z)

I'm doing the last hygiene pass now: remove the temporary coverage output, re-check the diff, and verify the branch base before staging only the intended P16 files.

369. assistant (2026-07-02T09:33:21.211Z)

The branch is still exactly on ``origin/master`` (``7b6544f``) before the P16 commit. One temporary coverage folder was under ``web/``, so I'm clearing that generated output before staging.

370. assistant (2026-07-02T09:33:32.308Z)

The working tree now contains only the intended P16 tracked edits plus the new resume E2E file; ignored build artifacts are staying ignored. I'm re-verifying the base one more time, then staging those exact paths.

371. assistant (2026-07-02T09:33:43.889Z)

Staging is cleanly scoped to the 19 intended P16 files. I'm committing now, then I'll push and open the PR with the coverage and screencast evidence in the body.

372. assistant (2026-07-02T09:33:49.016Z)

The commit is now a single clean branch commit (`7df0da5`) over `origin/master`. I'm checking GitHub auth, pushing it, and then I'll create the PR.

373. assistant (2026-07-02T09:34:02.495Z)

The branch is on GitHub. I'm opening this as a ready PR because your instruction is to merge the current fix once checks are good, and I'll include the local coverage gate plus the resume screencast evidence.

374. assistant (2026-07-02T09:34:18.697Z)

The GitHub app couldn't create the PR for this repo (`404` from the connector), so I'm using the authenticated `gh` CLI fallback on the already-pushed branch.

375. assistant (2026-07-02T09:34:38.735Z)

PR #144 is open: <https://github.com/avidullu/khelsutra/pull/144>. I'm watching the checks now; once they're green I'll merge it and then pause.

376. assistant (2026-07-02T09:34:46.551Z)

CI has started: `site` and `layout` are running, `web` is queued, and the Cloudflare Workers build already passed. I'll let the check run finish before merging.

377. assistant (2026-07-02T09:36:44.465Z)

All PR checks are green (`site`, `layout`, `web`, and the Workers build). I'm merging with the head SHA pinned so nothing else can slip into the merge.

378. assistant (2026-07-02T09:36:58.062Z)

The merge itself succeeded; the follow-up metadata query tripped on a wrong field name, so I'm re-querying the PR with GitHub's accepted fields and checking the local repo state.

379. assistant (2026-07-02T09:37:09.932Z)

P16 is merged as PR #144 and local `master` is at the merged commit. I'm updating the shared handoff/current-state notes now, while still leaving the engine tracker untouched per your latest constraint.

380. assistant (2026-07-02T09:37:37.254Z)

Before I write the handoff, I'm doing one final allowed-repo sync/check across `khelsutra`, `rally-annotator`, and `sports-data-collector`. I'm deliberately not touching `badminton-highlight-indexer`.

381. assistant (2026-07-02T09:37:57.750Z)

The allowed repos are synced: `khelsutra` is at `b5c7861`, `rally-annotator` at `ef9e088`, and `sports-data-collector` at `01ceda9` with only its pre-existing dirty `.gitignore` hunk. I'm recording the updated pending list in memory rather than editing the engine tracker.

382. assistant (2026-07-02T09:38:26.244Z)

The current-state note is updated with the P16 merge and the corrected pending list. I'm tightening the handoff top now so the next session lands on the newest truth instead of the

earlier engine pause.

383. assistant (2026-07-02T09:39:30.455Z)

P16 is merged and I'm pausing here as requested.

Merged: [khelsutra #144](https://github.com/avidullu/khelsutra/pull/144) at `b5c7861`. CI was green, and local coverage stayed above floors: statements `94.09%`, branches `87.64%`, functions `88.40%`, lines `94.09%`.

I updated the shared handoff/current-state notes with the new ground truth and pending list: [current-state.md](/C:/Users/avidu/OneDrive/Documents/claude-sync/memory/khelsutra-guru/current-state.md) and [session-handoff.md](/C:/Users/avidu/OneDrive/Documents/claude-sync/memory/khelsutra-guru/session-handoff.md). I did not edit `badminton-highlight-indexer`.

****Pending Audit Items****

ID	Scope	Status
P13	`sports-data-collector`	Pending owner decision: pre-existing dirty `.gitignore` artifact hunk preserved, not staged.
P17	`khelsutra`	Pending: show uploaded filename during upload/processing.
P18	`khelsutra`	Pending: hide/replace demo rally grid during real job.
P19	`khelsutra`	Pending: sanitize HTML/proxy 5xx API error bodies.
P20	`khelsutra`	Pending: Web Share API reel file-share attempt plus WhatsApp fallback.
P21	`khelsutra`	Pending: stale SPA API comment anchors.
P27	non-git/GCS	Partially blocked: metadata backup is owner-gated; engine doc update blocked by no-engine-edits.
P28	`sports-data-collector` / non-git	Blocked on owner/F: mount data access.

Engine-side rows P8-P10, P15, and P23-P26 remain excluded while the no-engine-edits constraint stands.

384. user (2026-07-02T12:16:21.679Z)

please go ahead and try to complete as many as possible by making and merging their PRs

385. assistant (2026-07-02T12:16:37.565Z)

I'll resume the allowed-repo track and keep `badminton-highlight-indexer` read-only. I'm going to start with the next `khelsutra` UI rows, one small PR at a time, and merge each after local verification plus green GitHub checks.

386. assistant (2026-07-02T12:16:47.663Z)

First I'm refreshing the allowed repos before reading anything in them, then I'll use the handoff/current-state notes as the source of truth for the pending list.

387. assistant (2026-07-02T12:16:56.545Z)

The memory has a newer engine-track checkpoint from another session, but your last repo-specific guard still says not to edit `badminton-highlight-indexer`, so I'm keeping this thread on the allowed repos. Next up is `khelsutra` P17, then I'll continue down the SPA rows if the tree stays clean.

388. assistant (2026-07-02T12:17:02.614Z)

I'm cutting the P17 branch directly from `origin/master` now, then I'll inspect the ingest components and tests around upload/processing status.

389. assistant (2026-07-02T12:17:15.133Z)

I'm reading the current upload/ingest path and the local test shape so this stays a tiny UI PR instead of growing sideways.

390. assistant (2026-07-02T12:17:37.676Z)

The simplest safe shape for P17 is panel-level: remember the user-facing source label when a file/URI is chosen, then render it in both the upload progress and engine polling states. I'm going to add one localized status line and pin it with component/upload tests.

391. assistant (2026-07-02T12:17:56.692Z)

I'm making the P17 edit now: a small source-name status line, cleared on cancel/reset, with locale keys mirrored through the site copy.

392. assistant (2026-07-02T12:18:17.294Z)

The patch missed one JSX context, so I'm reading the nearby lines and applying it in smaller chunks.

393. assistant (2026-07-02T12:18:22.605Z)

I've got the exact JSX now. The edit is still small: helper + remembered source name + status-line render in upload/store/process states.

394. assistant (2026-07-02T12:18:52.573Z)

The component edit is in. I'm adding the matching i18n keys across the six catalogs, then I'll regenerate the site locale mirror and run focused tests.

395. assistant (2026-07-02T12:19:07.591Z)

Some translated strings differ from my guess, so I'm locating the existing `uploadingStatus` keys and patching against the actual catalog text.

396. assistant (2026-07-02T12:19:10.892Z)

PowerShell didn't like the POSIX-style glob there; I'm rerunning the search with `rg`'s own file filter.

397. assistant (2026-07-02T12:19:15.359Z)

I found the actual catalog text. I'll patch the root locale files and then let the existing copy script update `site/public/locales`.

398. assistant (2026-07-02T12:19:27.828Z)

I'm pinning the behavior in the existing upload integration test: it should show the chosen filename immediately and keep showing it once the job has moved into processing.

399. assistant (2026-07-02T12:19:34.510Z)

The UI test is updated. I'm mirroring the locale catalogs into `site/public` with the repo script, then running the focused web/site checks for this PR.

400. assistant (2026-07-02T12:19:44.384Z)

Because this touches the SPA UI, I'll include the normal web coverage/type gates and the Playwright screencast run. The focused tests run first so we catch any local issue quickly.

401. assistant (2026-07-02T12:19:57.854Z)

Focused web tests passed (`68` tests). The site test command failed only because I gave Node's test runner a Jest flag; I'm rerunning it in the repo's actual form, then I'll run the full web gates.

402. assistant (2026-07-02T12:20:08.627Z)

Site tests and web typecheck are green. I'm moving to the coverage/build gates plus the existing Playwright screencast suite for UI evidence.

403. assistant (2026-07-02T12:21:16.720Z)

Coverage, build, and the site browser checks are green. I'm reading the coverage summary directly, then I'll run the Playwright screencast suite and inspect the generated evidence path.

404. assistant (2026-07-02T12:21:20.361Z)

The coverage gate passed but this repo doesn't leave `coverage-summary.json` at that path. I'm locating the generated coverage artifact and then I'll capture the screencast evidence.

405. assistant (2026-07-02T12:21:27.763Z)

No summary file was left behind, so I'll rerun coverage once with a JSON/text summary reporter. It's the same coverage gate, just with a durable number to put in the PR.

406. assistant (2026-07-02T12:22:09.297Z)

The coverage gate is above the repo floor, but the exact total dipped versus the previous checkpoint because the new UI branches added lines. I'm adding a little more behavioral coverage so this PR doesn't carry a coverage regression.

407. assistant (2026-07-02T12:22:21.047Z)

I only generated the JSON summary, not the per-line coverage file. I'm creating the per-file map once so I can target the uncovered new branches instead of guessing.

408. assistant (2026-07-02T12:22:38.357Z)

I found the new uncovered area: it's mostly the fallback/source-name cases and cancel/reset clearing. I'm adding two tiny UI tests to cover those actual behaviors, then I'll rerun full coverage.

409. assistant (2026-07-02T12:22:52.752Z)

I'm adding two coverage-bearing assertions: one covers a pasted host-path basename and cancel clearing it, the other covers retry/reset after an error. That should bring the new UI lines back into the green zone without inventing irrelevant tests.

410. assistant (2026-07-02T12:23:06.121Z)

Those extra assertions are in. I'm rerunning the focused panel/upload tests, then the full coverage summary again to confirm the numbers recovered.

411. assistant (2026-07-02T12:23:14.298Z)

Focused tests are green. I'm rerunning the full coverage gate now and will use its JSON summary as the PR coverage record.

412. assistant (2026-07-02T12:23:28.745Z)

Full coverage passed again. I'm grabbing the exact totals and then running Playwright for the UI artifact.

413. assistant (2026-07-02T12:24:37.927Z)

I'll pick up the in-flight `khelsutra` P17 branch first, then keep moving through the allowed repos only. I'm going to re-check the branch, HEAD, and local verification state before staging anything so the shared checkout stays tidy.

414. assistant (2026-07-02T12:24:45.343Z)

I'm using the GitHub publish workflow for the PR plumbing. Local `git` will do the branch/commit/push work; I'll use `gh` where the connector can't handle this repo cleanly.

415. assistant (2026-07-02T12:25:00.472Z)

The P17 branch is still on the expected base and the tree only contains the ingest filename visibility files. I'm going to run the remaining UI evidence check now, then I'll do one final coverage pass and commit only those explicit paths.

416. assistant (2026-07-02T12:25:13.443Z)

The web package has a dedicated `screenshot` script that builds and runs the Playwright suite with video retention. I'm checking the config only enough to know where evidence lands, then I'll run it.

417. assistant (2026-07-02T12:25:23.357Z)

One thing I'm tightening before the PR: the statement/function coverage is up, but branch coverage is a hair under the previous local number. I'm going to add a very small assertion that exercises another UI branch, so the coverage story is clean rather than technically arguable.

418. assistant (2026-07-02T12:25:45.794Z)

The existing coverage report was JSON-summary only, so it doesn't show missed branch locations. I'm going to add coverage where it matters functionally here: a URL basename with decoding, plus the non-local store path keeping the original upload filename visible.

419. assistant (2026-07-02T12:26:02.483Z)

Those extra assertions are in. I'm rerunning the focused ingest tests first, then the full coverage command so we get fresh totals for the PR body.

420. assistant (2026-07-02T12:26:13.818Z)

Focused tests are green. Now I'm running the full web coverage report and the Playwright screenshot suite; this is the longer local gate for the UI PR.

421. assistant (2026-07-02T12:27:23.214Z)

Playwright passed and generated fresh raw video under the shared screenshot artifact folder. I'm doing the remaining repo checks now: typecheck, site test suite, and a whitespace diff check before staging.

422. assistant (2026-07-02T12:27:39.563Z)

Coverage is still comfortably above the repo floor, but branch percentage is just under the prior local P16 note by 0.02. I'm going to inspect the detailed coverage map and add one focused test rather than hand-waving that away.

423. assistant (2026-07-02T12:28:14.363Z)

The detailed coverage map pointed at an uncovered medium-picker branch. I'm adding a tiny Drive-medium test; it's useful coverage, not just a number-padder.

424. assistant (2026-07-02T12:28:26.590Z)

Drive-medium coverage is in. I'll rerun the focused panel tests and then the full coverage report one more time for final PR numbers.

425. assistant (2026-07-02T12:28:53.015Z)

Coverage is now clean: 94.65% statements/lines, 89.92% functions, 87.92% branches. I'm doing a final status/branch verification and collecting the screenshot artifact path before staging the

explicit files.

426. assistant (2026-07-02T12:29:02.641Z)

The final local checks are green. I'm fetching now and re-verifying that this branch is still based where expected before I stage only the intended files.

427. assistant (2026-07-02T12:29:15.432Z)

Origin/master still matches the branch base. I'm staging only the locale mirrors and ingest component/test files for P17 now.

428. assistant (2026-07-02T12:29:30.031Z)

The staged set matches the intended PR exactly. I'm committing this as the P17 fix, then I'll push and open the PR with the verification notes and screencast path.

429. assistant (2026-07-02T12:29:40.062Z)

Commit `5e00b90` is the only commit over `origin/master`, so the branch is cleanly scoped. Pushing now.

430. assistant (2026-07-02T12:29:49.631Z)

The branch is pushed. I'm opening the PR with the verification log and then I'll watch checks before merging it.

431. assistant (2026-07-02T12:30:31.518Z)

PR #145 is open. I'm checking its CI state now and will merge with the exact head SHA once it's green.

432. assistant (2026-07-02T12:32:25.208Z)

All three CI checks for #145 are green. I'm merging it with the pinned head SHA now, then I'll fast-forward local master before cutting the next branch.

433. assistant (2026-07-02T12:32:31.703Z)

The GitHub connector can't see this repo for merges, so I'm falling back to `gh pr merge` with the same pinned head SHA.

434. assistant (2026-07-02T12:32:53.044Z)

PR #145 is merged and the local checkout fast-forwarded to the squash commit `07b935d`. I'm refreshing the handoff/current-state now, then I'll cut the next allowed-repo branch from `origin/master` for P18.

435. assistant (2026-07-02T12:32:59.132Z)

Memory currently still says ?pause after P16,? so I'm correcting that to the real shipped state: P17/#145 merged, next is P18, and P17 is removed from the pending allowed-repo list.

436. assistant (2026-07-02T12:34:08.871Z)

Handoff/current-state now reflect P17 merged and P18 next. I'm pulling `khesutra` master, then I'll branch from `origin/master` for the demo-grid-while-ingesting fix.

437. assistant (2026-07-02T12:34:24.200Z)

The P18 branch is cut from remote master. I'm reading the SPA ingest/App flow now to find the smallest existing-pattern hook for ?real job active, don't show demo rallies.?

438. assistant (2026-07-02T12:34:34.647Z)

App currently seeds demo rallies until `videoId` exists, while `IngestPanel` owns job busy state internally. The fix should expose a tiny busy callback from `IngestPanel` and let `App` replace the pick/build demo surface while that callback says a real ingest is active.

439. assistant (2026-07-02T12:35:01.399Z)

I've got the shape: no new global state machinery, just one `onBusyChange` callback from the existing ingest component. While that callback is true and no real `videoId` exists yet, `App` will show the loading banner and suppress the demo query/grid/footer surface.

440. assistant (2026-07-02T12:35:45.759Z)

The P18 code and a focused App test are in. I'm running the new test with the related ingest panel tests first so we can catch any state-timing weirdness before the wider suite.

441. assistant (2026-07-02T12:36:02.681Z)

The implementation looks right; the new test is what timed out because it used fake timers with Testing Library's `waitFor`. I'm switching that test to real timers and letting the promises settle normally.

442. assistant (2026-07-02T12:36:08.974Z)

Test timing is patched. I'm rerunning the focused App/ingest set now.

443. assistant (2026-07-02T12:36:22.992Z)

The focused tests are green. I'm also pinning P18 in the Playwright ingest CUJ, so the screencast/CI path explicitly proves the demo grid disappears during a live ingest.

444. assistant (2026-07-02T12:36:37.263Z)

P18 is now covered in unit and Playwright. I'm running the main local gate: typecheck, full web coverage, and the screencast suite.

445. assistant (2026-07-02T12:39:08.124Z)

I'll pick up from the in-progress `khelsutra` P18 branch, first doing the required freshness check. Since this branch may have local uncommitted work, I'll treat any pull refusal as a safety signal and avoid touching unrelated state.

446. assistant (2026-07-02T12:39:13.510Z)

Freshness check passed for `khelsutra`; the branch is already up to date. I'm going to use the GitHub publish workflow guidance for the PR/merge steps, then finish the local verification gate before staging.

447. assistant (2026-07-02T12:39:30.515Z)

The P18 code path already has the focused tests, typecheck, full web coverage, and screencast run from the previous stretch. I'm filling the remaining local gate now: site tests, whitespace diff check, exact coverage totals, and the latest screencast evidence paths.

448. assistant (2026-07-02T12:39:40.520Z)

All P18 local gates are green so far: site tests passed, whitespace check is clean, and web coverage is above the floor. I'm doing the required branch/HEAD re-check now, then I'll stage only the four intended files.

449. assistant (2026-07-02T12:39:54.038Z)

The branch is exactly where it should be: `fix/hide-demo-rallies-while-ingesting` at the same HEAD as `origin/master`, with no extra commits. I'm staging just the P18 files and committing them now.

450. assistant (2026-07-02T12:40:00.676Z)

Staging is explicit and cleanly scoped. I'm checking the staged diff/stat once before commit so no extra file slips into the PR.

451. assistant (2026-07-02T12:40:03.087Z)

The staged set is exactly the four intended files. Commit next, then I'll push and open the PR.

452. assistant (2026-07-02T12:40:11.434Z)

P18 is committed locally. I'm doing the post-commit branch audit now: the PR branch should contain only this one commit over `origin/master`.

453. assistant (2026-07-02T12:40:14.418Z)

The branch contains only the P18 commit over `origin/master`, and the worktree is clean. I'm pushing it and opening the PR with the verification notes included.

454. assistant (2026-07-02T12:40:25.003Z)

Before I publish the PR, I'm checking whether this repo has a web lint script I should run. Typecheck, tests, coverage, and screencast are already covered; lint should be included if the package exposes it.

455. assistant (2026-07-02T12:40:36.251Z)

There's no web lint script in `package.json`, so I'll record that plainly. I'm also checking the screencast artifact directory for the ingest recording name so the PR evidence points at the most relevant video, not just the newest generic runs.

456. assistant (2026-07-02T12:40:50.417Z)

The suite passed, but the artifact copier didn't leave a distinct `ingest` video in `\_raw`; I'm checking the local Playwright/screencast wiring before deciding whether to regenerate targeted evidence for this exact busy-ingest state.

457. assistant (2026-07-02T12:41:01.263Z)

The screencast config does record video locally outside git, so I'm going to rerun just the ingest Playwright case to produce a focused artifact for the PR. This is short and gives us evidence for the exact UI state this change touches.

458. assistant (2026-07-02T12:41:29.616Z)

The targeted Playwright run passed, but it cleaned `\_raw` and did not leave a retained video for that single-case invocation. I'm going to rerun the repo's normal `npm run screencast` command to restore convention-shaped evidence before I publish the PR.

459. assistant (2026-07-02T12:42:27.441Z)

The full screencast command restores the standard artifact set, but the ingest case still doesn't leave a separate video even though the test itself passes. I'll cite the passing ingest Playwright check plus the retained raw suite videos, and I'll keep the PR body honest about that artifact shape.

460. assistant (2026-07-02T12:42:36.833Z)

To make the UI evidence more directly useful, I'm going to add a non-git screenshot artifact

under `khelsutra-screenshot/\_raw` for the active-ingest replacement state. That avoids changing the repo while giving the PR reviewer a precise visual proof point.

461. assistant (2026-07-02T12:42:49.485Z)

For the screenshot I'll use the same built SPA plus `e2e/mock-engine.mjs`, drive the URI ingest flow to the loading state, and save a PNG under the non-git `\_raw` artifact directory.

462. assistant (2026-07-02T12:43:01.935Z)

The ingest Playwright file intentionally disables video (`test.use({ video: "off" })`), which explains the missing video artifact. I'm saving a targeted screenshot instead, while still keeping the passing ingest test as the behavioral proof.

463. assistant (2026-07-02T12:43:22.192Z)

The screenshot artifact was created successfully. I'm visually checking it once so the PR evidence is not just "a file exists," but actually captures the intended state.

464. assistant (2026-07-02T12:43:27.513Z)

The screenshot shows the real ingest panel active with the sample rally grid and footer gone, which is the exact P18 behavior. I'm checking git status once more, then I'll create the PR.

465. assistant (2026-07-02T12:43:33.109Z)

The repo worktree is clean. I'm creating the P18 PR now with the coverage numbers and the focused screenshot path in the body.

466. assistant (2026-07-02T12:43:51.988Z)

PR #146 is open. I'm watching the GitHub checks now; if they go green, I'll merge this one before moving to the next allowed audit item.

467. assistant (2026-07-02T12:46:03.822Z)

GitHub checks for #146 are green (`layout`, `site`, `web`). I'm grabbing the PR head SHA and merging with `--match-head-commit` so the merge is pinned to the checked commit.

468. assistant (2026-07-02T12:46:08.817Z)

The checked head is `1ec4e6f1eebbb5fdd298f631c18cdf9d41a72055` and the PR is clean to merge. Merging P18 now.

469. assistant (2026-07-02T12:46:34.293Z)

P18 is merged as `a7ae666` on `master`. I'm updating the local handoff/current-state tracker now, then I'll continue with the next allowed non-engine item.

470. assistant (2026-07-02T12:46:43.003Z)

I'm applying the tracker update in the OneDrive memory files now: P18 becomes done, `khelsutra` head advances to `a7ae666`, and the pending allowed list starts at P19/P20/P21 plus the gated cleanup rows.

471. assistant (2026-07-02T12:47:24.909Z)

The handoff is current for P18. I'm applying the same checkpoint to `current-state.md` so the short tracker view agrees with the longer handoff.

472. assistant (2026-07-02T12:47:55.892Z)

Tracker state is updated. Next I'm moving to P19 in `khelsutra`: sanitize ugly/proxy HTML 5xx

bodies before they show up in SPA API errors. I'll pull master first, then branch fresh from `origin/master`.

473. assistant (2026-07-02T12:48:01.365Z)

`khelsutra` master is fresh. I'm creating the P19 branch explicitly from `origin/master` now, per the shared-checkout safety rule.

474. assistant (2026-07-02T12:48:16.089Z)

Branch `fix/sanitize-api-error-bodies` is cut from `origin/master`. I'm reading the SPA API/error plumbing and the existing tests to make the smallest P19 fix in the right place.

475. assistant (2026-07-02T12:48:24.997Z)

The right layer is `web/src/lib/errors.ts`: it already distinguishes network-ish 5xx from engine details, but its parser can still treat HTML/proxy pages as user-facing detail. I'm narrowing that parser and keeping real JSON `{detail}` messages intact.

476. assistant (2026-07-02T12:48:43.405Z)

The fix will stay tiny: keep real FastAPI JSON details and plain engine messages, but suppress browser/proxy HTML boilerplate and known proxy text before classification. Then a 5xx with only that suppressed body falls back to the existing friendly network message.

477. assistant (2026-07-02T12:48:54.451Z)

The parser now filters HTML/proxy boilerplate. I'm adding regression tests so future 5xx proxy pages can't leak into the ingest/build UI as raw markup.

478. assistant (2026-07-02T12:49:11.009Z)

The regression tests are in place. I'm running the focused error/client/build tests first; if they're clean, I'll run the broader web gate and coverage.

479. assistant (2026-07-02T12:49:32.473Z)

Focused tests passed: 42 tests across parser, error text, client, and build-error UI. I'm running the full web coverage/typecheck/site gate next, and I'll also produce a small screenshot artifact for the sanitized user-facing error state.

480. assistant (2026-07-02T12:49:59.551Z)

The main local gates are green so far: typecheck, full web coverage, site tests, and whitespace. I'm running the normal Playwright screencast suite and then I'll create a focused screenshot showing the sanitized proxy-error UI.

481. assistant (2026-07-02T12:50:55.983Z)

The standard screencast suite is green again. For the P19-specific visual, I'm going to serve the built SPA, mock a 502 HTML body from `/api/compile`, and capture the friendly error instead of raw markup.

482. assistant (2026-07-02T12:51:06.393Z)

I've got the exact interaction surface now: load `?video=vidX`, let the mocked rallies populate, click Build, and return a 502 HTML page from `/api/compile`. The screenshot should show the friendly "can't reach" line, with no `<html>`/Bad Gateway text.

483. assistant (2026-07-02T12:51:32.399Z)

The P19 screenshot was generated and the script asserted that no proxy boilerplate leaked into the visible page. I'm visually checking the artifact once, then I'll prepare the PR.

484. assistant (2026-07-02T12:51:37.580Z)

The screenshot looks right: the user sees the friendly network message, not the 502 HTML. I'm doing one small readability touch in the predicate, then I'll re-run the focused parser test and prep the branch for commit.

485. assistant (2026-07-02T12:51:49.163Z)

That's just a parenthesized condition, but I'm still rerunning the tight parser test and collecting the exact coverage totals before the branch audit.

486. assistant (2026-07-02T12:52:00.485Z)

P19 checks are still green after the readability touch, with web coverage at 94.67% lines/statements. I'm re-fetching and re-checking branch/HEAD now, then I'll stage only the two intended files.

487. assistant (2026-07-02T12:52:18.203Z)

The branch is still exactly on `origin/master` (`a7ae666`) with no extra commits, and the diff check is clean. I'm staging just `errors.ts` and its test.

488. assistant (2026-07-02T12:52:25.433Z)

Staged scope is explicit. I'm checking the cached diff and file list one last time before commit.

489. assistant (2026-07-02T12:52:28.959Z)

The cached diff is only the intended parser and regression tests. I'm committing P19 now.

490. assistant (2026-07-02T12:52:38.980Z)

P19 is committed locally. I'm confirming the branch contains only this commit over `origin/master`, then I'll push and open the PR.

491. assistant (2026-07-02T12:52:43.715Z)

The branch has only the P19 commit and the repo tree is clean. Pushing and opening PR # next.

492. assistant (2026-07-02T12:52:51.071Z)

The P19 branch is on GitHub. I'm creating the PR with the local verification, coverage totals, and the screenshot evidence path.

493. assistant (2026-07-02T12:53:16.773Z)

PR #147 is open. I'm watching the GitHub checks; if they're green, I'll merge with a pinned head SHA.

494. assistant (2026-07-02T12:55:29.597Z)

GitHub checks for #147 are green (`layout`, `site`, `web`). I'm fetching the checked head SHA and merging with `--match-head-commit`.

495. assistant (2026-07-02T12:55:37.088Z)

The checked head is `caacae06ec52043b50f04a116afd12830abab0e9` and the PR is clean. Merging P19 now.

496. assistant (2026-07-02T12:55:54.553Z)

P19 is merged as `2daa0ad` on `master`. I'm updating the tracker memory, then I'll continue to

P20.

497. assistant (2026-07-02T12:56:03.095Z)

Memory still shows P19 pending, so I'm moving it to completed now and advancing the allowed `khelsutra` head to `2daa0ad`.

498. assistant (2026-07-02T12:56:40.421Z)

The long handoff now reflects P19 as merged. I'm updating the short `current-state` view to match before starting P20.

499. assistant (2026-07-02T12:57:12.998Z)

Current-state is aligned. P20 is next: try native Web Share file sharing for reels, with WhatsApp/link fallback. I'll pull master again before reading and branch fresh from `origin/master`.

500. assistant (2026-07-02T12:57:18.761Z)

`khelsutra` master is fresh at `2daa0ad`. I'm cutting the P20 branch explicitly from `origin/master`.

501. assistant (2026-07-02T12:57:36.839Z)

The P20 branch is ready. I'm reading the reel result component and existing tests/locales so the share control fits the current UI instead of bolting on a new pattern.

502. user (2026-07-02T12:57:38.006Z)

Make sure to update the tracker to reflect all that is completed

503. assistant (2026-07-02T12:58:00.065Z)

Already on it: since the engine repo remains off-limits in this allowed-repos track, I'm treating the OneDrive handoff/current-state files as the live tracker mirror. I've marked P18 and P19 complete there with PR numbers, merge SHAs, coverage, checks, and evidence paths; I'm doing one quick readback now so the tracker and pending list are definitely aligned before I continue P20.

504. assistant (2026-07-02T12:58:14.452Z)

Tracker readback confirms the completed rows are up to date: P18 via #146 and P19 via #147 are no longer pending, and the remaining allowed list is P13, P20, P21, P27, P28. I'm continuing with P20 now on the fresh branch.

505. assistant (2026-07-02T12:59:31.557Z)

I'll keep the tracker current as I go. I'm resyncing the allowed checkout state first, then I'll continue from the in-progress `khelsutra` P20 branch and avoid the engine repo entirely.

506. assistant (2026-07-02T12:59:36.766Z)

The workspace root is just a container folder, not a git repo, so there's nothing to fast-forward there. I'm applying the freshness rule at the actual sub-repo level now.

507. assistant (2026-07-02T12:59:48.529Z)

`khelsutra` is fast-forward clean from its upstream. I'm checking the branch/HEAD and reading the live handoff/current-state so the pending list and tracker edits line up with what's already shipped.

508. user (2026-07-02T12:59:49.264Z)

sounds great! Please carry on and run through these items. One note is not not work on the main.py refactoring. Also make sure sync to remote head

509. assistant (2026-07-02T13:00:01.604Z)

Using the GitHub publishing skill for the PR/merge workflow, and I'll keep the engine/main.py split off-limits. I'll resync `khelsutra` to remote head again before touching the P20 branch.

510. assistant (2026-07-02T13:00:10.980Z)

The publishing workflow is loaded. I'm going to keep this as a small `khelsutra` PR: native file share attempt for the finished reel, while leaving the existing WhatsApp link as the fallback path.

511. assistant (2026-07-02T13:00:28.892Z)

Remote head is `2daa0ad`, and the P20 branch is exactly on it with a clean tree. I'm reading the component, tests, and locale/package wiring before editing so the new share path fits the existing UI and coverage shape.

512. assistant (2026-07-02T13:00:37.179Z)

The existing result card already has the right fallback anchor and localized copy. I'm going to add a separate native-share button with success/failure status text, then mirror those new locale keys into the site's public locale bundle.

513. assistant (2026-07-02T13:01:17.276Z)

I'm about to edit `CompileResult` and the result locale keys only. The component change will keep the downloaded file path untouched, fetch the MP4 into a `File`, ask `navigator.canShare` before sharing, and fall back to the existing WhatsApp path if the browser can't attach the file.

514. assistant (2026-07-02T13:01:41.800Z)

The component is patched. Next I'm adding the localized labels/status strings in the root catalogs, then I'll run the repo's locale copy script so the marketing/static bundle stays byte-identical.

515. assistant (2026-07-02T13:02:07.442Z)

One locale hunk didn't match exactly, so the multi-file patch didn't apply. I'm reading the compact result blocks and will patch the catalogs against their actual text.

516. assistant (2026-07-02T13:02:11.676Z)

The catalog text varies slightly by locale, which is good news for real translations and bad news for lazy patching. I've got the exact blocks now and am adding the four new result strings per language.

517. assistant (2026-07-02T13:02:33.935Z)

Root locale catalogs are updated. I'm mirroring them through the project's own `copy-locales` script now, which keeps the static site drift guard happy without hand-editing duplicate JSON.

518. assistant (2026-07-02T13:02:38.931Z)

The duplicate locale bundle is synced. I'm adding focused component tests now: one proves we fetch and attach the MP4 to Web Share, and one proves a non-shareable file leaves users on the WhatsApp/download fallback.

519. assistant (2026-07-02T13:03:02.121Z)

One assertion used a matcher this repo doesn't appear to install. I'm swapping it to the project's existing plain `textContent` style before running tests.

520. assistant (2026-07-02T13:03:13.266Z)

The scoped diff is only the result card, its focused tests, and locale mirrors. I'm running the focused Vitest slice first; if that's happy I'll widen out to typecheck, full coverage, screencast, and site tests.

521. assistant (2026-07-02T13:03:23.050Z)

Focused tests are green: 8 passing. I'm moving to TypeScript now; this is where the browser Web Share typings will complain if the shape is off.

522. assistant (2026-07-02T13:03:30.670Z)

TypeScript caught only the test mock-call typing, not the component. I'm tightening that assertion cast rather than weakening the production code.

523. assistant (2026-07-02T13:03:37.254Z)

Test typing is cleaned up. I'm rerunning the focused test and typecheck together so we don't carry a stale green.

524. assistant (2026-07-02T13:03:45.196Z)

Focused test and typecheck are both green. I'm running full web coverage next so we can explicitly verify it stays above the repo floor and doesn't dip from the prior checkpoint.

525. assistant (2026-07-02T13:04:00.618Z)

Full web coverage passed. I'm pulling the coverage JSON totals now, then I'll run the UI screencast suite and the static site tests.

526. assistant (2026-07-02T13:04:08.835Z)

Coverage is still far above the floor, but it did dip versus the prior checkpoint (`94.41/90.14/87.5` vs `94.67/90/88.11`). Since you asked for no regression, I'm adding a couple of focused cases to cover the fallback/error branches instead of accepting that slip.

527. assistant (2026-07-02T13:04:17.815Z)

The uncovered bits are all inside `CompileResult`, mostly the defensive native-share branches. I'm going to cover unsupported browser, failed fetch, abort-cancel, and maybe non-MP4 blob fallback so the aggregate moves back up.

528. assistant (2026-07-02T13:04:49.905Z)

A small refinement: the `?unsupported before click?` guard is unreachable because the native-share button is only rendered when support exists. I'm removing that dead defensive path and adding tests for share failure and user-cancel paths, which is cleaner and should pull coverage back up.

529. assistant (2026-07-02T13:05:16.878Z)

The dead path is gone and the extra share-failure tests are in. I'm rerunning the focused slice plus typecheck before recalculating full coverage.

530. assistant (2026-07-02T13:05:24.721Z)

Focused tests are now 10 passing and typecheck is still clean. Re-running full coverage to confirm the no-regression target before I move on.

531. assistant (2026-07-02T13:05:49.968Z)

Coverage is now back above the prior checkpoint: `94.78` lines/statements, `90.14` functions, `88.22` branches. I'm running the UI screencast suite and static-site tests next; after that I'll grab a P20-specific screenshot artifact.

532. assistant (2026-07-02T13:06:47.707Z)

Screencast suite is green (5 Playwright tests) and site tests are green (95). I'm creating a P20-specific screenshot under the existing `\_raw` convention so the PR has visual evidence of the new native share button/fallback surface.

533. assistant (2026-07-02T13:07:11.431Z)

The generic screencast videos exist. For the P20-specific evidence, I'm using the same mock-engine harness on a throwaway port, driving the real built SPA to a completed reel, clicking the native share button with `canShare=false`, and saving the visible fallback state as a PNG under `\_raw`.

534. assistant (2026-07-02T13:07:32.512Z)

The P20 screenshot artifact is saved. I'm opening it once to make sure it actually shows the completed reel, Share video control, WhatsApp fallback, and fallback status text cleanly.

535. assistant (2026-07-02T13:07:45.759Z)

The screenshot looks right: completed reel, download, native `Share video`, WhatsApp fallback, and the fallback status line are all visible without overlap. I'm doing whitespace/status checks now, then I'll stage only the intended files.

536. assistant (2026-07-02T13:07:59.200Z)

The diff is clean and scoped. Before staging, I'm doing the required branch/HEAD recheck against `origin/master` so no sibling work sneaks into this PR.

537. assistant (2026-07-02T13:08:04.410Z)

Branch and `origin/master` both point at `2daa0ad`, and there are no commits above base yet. I'm staging only the fourteen intended paths now.

538. assistant (2026-07-02T13:08:17.174Z)

Only the intended result-card/test/locale paths are staged, and the cached diff check is clean. I'm committing P20 now.

539. assistant (2026-07-02T13:08:28.806Z)

Commit `897f54a` is created. I'm doing the post-commit base-range check, then I'll push and open the PR with the verification and screenshot evidence included.

540. assistant (2026-07-02T13:08:31.750Z)

The branch contains exactly one commit above `origin/master`. Pushing it now.

541. assistant (2026-07-02T13:08:43.953Z)

The branch is on GitHub. I'm opening the PR now with the local verification, coverage totals, and the P20 screenshot path in the body so CI and the tracker have the same ground truth.

542. assistant (2026-07-02T13:08:59.842Z)

PR #148 is open. I'm watching the GitHub checks now; if they go green, I'll squash-merge it and immediately update the tracker/handoff to mark P20 complete.

543. assistant (2026-07-02T13:11:04.691Z)

PR #148 checks are all green: `site`, `layout`, and `web`. I'm squash-merging it now and deleting the remote branch.

544. assistant (2026-07-02T13:11:19.658Z)

#148 is merged as `816a088`. I'm checking the local branch state, then I'll refresh the memory tracker to move P20 from pending to complete.

545. assistant (2026-07-02T13:11:29.303Z)

The merge is confirmed: `master` and `origin/master` are both at `816a088`. I'm updating the handoff/current-state tracker now so P20 is recorded as shipped and the pending list moves on to P21.

546. assistant (2026-07-02T13:12:42.663Z)

Tracker memory is patched to mark P20 complete and P21 next. I'm reading back the changed lines so I don't accidentally leave the handoff in an aspirational state.

547. assistant (2026-07-02T13:12:55.320Z)

P20 is fully merged and the tracker readback shows it as complete. I'm moving to P21 next, staying in `khelsutra`, explicitly syncing `master`, then branching from `origin/master`; no engine/main.py work.

548. assistant (2026-07-02T13:13:01.991Z)

`khelsutra` is already at remote head `816a088`. I'm cutting the P21 branch directly from `origin/master` now, per the shared-checkout rule.

549. assistant (2026-07-02T13:13:16.196Z)

The P21 branch is fresh from `origin/master`. I'm searching for stale backend/API line anchors in the SPA and adjacent docs/comments now.

550. assistant (2026-07-02T13:13:38.769Z)

The search found the intended stale anchors in SPA code comments: `web/src/api/types.ts`, `web/src/lib/floor.ts`, and `web/src/components/CompileResult.tsx` still cite moving `backend/main.py` line numbers. I'm keeping P21 scoped to those SPA comments and changing them to symbol/route-based anchors.

551. assistant (2026-07-02T13:13:43.383Z)

The first pass exposed a few more stale line anchors in `web/src/api/client.ts`, so I'm including those in the same P21 comment-only cleanup. No behavior changes.

552. assistant (2026-07-02T13:13:53.281Z)

I'm patching only comments: replacing drifting `main.py:<line>` / `uploads.py:<line>` citations with stable route or symbol anchors. This keeps the code behavior byte-for-byte the same.

553. assistant (2026-07-02T13:14:28.197Z)

The SPA code comments are patched. I'm running a targeted grep to confirm the stale `main.py:<line>` anchors are gone from `web/src`, while leaving legitimate examples like `localhost:8000` alone.

554. assistant (2026-07-02T13:14:38.846Z)

The stale anchor grep is clean for `web/src`. Since this touches SPA files, I'm still running

the web checks and screencast evidence, even though it's comment-only.

555. assistant (2026-07-02T13:15:03.789Z)

Typecheck, full web coverage, and site tests are green. I'm reading the coverage totals now, then I'll run the screencast suite for the UI evidence requirement.

556. assistant (2026-07-02T13:15:08.310Z)

Coverage is unchanged/improved at `94.78/90.14/88.25`, and the scoped diff check is clean. I'm running `npm run screencast` now.

557. assistant (2026-07-02T13:16:10.629Z)

The screencast suite passed again; the Vite chunk-size warning is the existing one. I'm doing the required branch/HEAD recheck before staging the comment-only P21 diff.

558. assistant (2026-07-02T13:16:16.018Z)

Branch and `origin/master` both remain at `816a088`, with no local commits above base. I'm staging only the four comment files for P21.

559. assistant (2026-07-02T13:16:25.773Z)

The four files are staged. I'm checking the cached diff and whitespace before committing.

560. assistant (2026-07-02T13:16:29.328Z)

Cached diff is clean, and the stale-anchor grep returns no matches. Committing P21 now.

561. assistant (2026-07-02T13:16:38.229Z)

Commit `40c0ca2` is ready. I'm verifying the branch contains exactly that one commit above `origin/master`, then I'll push and open #149.

562. assistant (2026-07-02T13:16:41.385Z)

The branch is a single P21 commit above `origin/master`. Pushing it now.

563. assistant (2026-07-02T13:16:48.256Z)

The P21 branch is up. I'm opening the PR with the grep result, coverage totals, and screencast evidence noted.

564. assistant (2026-07-02T13:17:02.720Z)

PR #149 is open. Watching CI now; when `site`, `layout`, and `web` are green I'll merge and update the tracker again.

565. assistant (2026-07-02T13:19:28.152Z)

PR #149 checks are green across `site`, `layout`, and `web`. Squash-merging it now.

566. assistant (2026-07-02T13:19:46.899Z)

#149 is merged as `61e5ccb`. I'm updating the tracker again to mark P21 complete; after that the remaining allowed rows are all gated unless a new allowed item has appeared in the sibling repos.

567. assistant (2026-07-02T13:19:52.994Z)

The merge is confirmed and local `master` is at `61e5ccb`. I'm patching the memory tracker now: P21 complete, pending list reduced to P13/P27/P28 with their gates.

568. assistant (2026-07-02T13:21:18.806Z)

The tracker is patched for #149. I'm doing one final allowed-repo sync pass for `rally-annotator` and `sports-data-collector` so the pending/gated state is based on current heads, not stale memory.

569. assistant (2026-07-02T13:21:33.742Z)

`rally-annotator` and `sports-data-collector` both fast-forward checked cleanly; no remote head moved. I'm checking their exact status/HEAD now before I wrap this pass.

570. assistant (2026-07-02T13:22:28.793Z)

Done. I resynced to remote head and carried through the remaining actionable allowed items.

Merged PRs:

- `khelsutra` #148: native reel file sharing fallback, with WhatsApp fallback preserved.
- `khelsutra` #149: replaced stale SPA API line-anchor comments, without touching `main.py` refactoring.

Verification stayed green:

- `khelsutra/web` typecheck clean.
- `khelsutra/web` coverage: `94.78%` statements/lines, `90.14%` functions, `88.25%` branches, no regression from the prior checkpoint.
- `khelsutra/web` screencast passed.
- `khelsutra/site` tests passed: `95 passed`.
- GitHub checks for both PRs passed before merge.

Tracker updated:

- `C:\Users\avidu\OneDrive\Documents\claude-sync\memory\khelsutra-guru\session-handoff.md`
- `C:\Users\avidu\OneDrive\Documents\claude-sync\memory\khelsutra-guru\current-state.md`

Current remaining allowed audit items are gated rather than directly actionable:

- `P13`: `sports-data-collector` `.gitignore` artifact policy needs owner decision; existing dirty `.gitignore` hunk preserved.
- `P27`: non-git/GCS metadata backup portion is owner-gated; engine doc side is excluded by no-engine-edits.
- `P28`: sports-data/non-git cleanup blocked on owner/F: mount availability.

I did not work in `badminton-highlight-indexer`, did not do the `main.py` refactor, and left `sports-data-collector`'s pre-existing `.gitignore` change untouched.

571. user (2026-07-02T13:56:05.314Z)

did you also update the tracker document and can you highlight what all items are left overs?

572. assistant (2026-07-02T13:56:23.079Z)

Yes. I updated the active memory tracker files:

- `C:\Users\avidu\OneDrive\Documents\claude-sync\memory\khelsutra-guru\current-state.md`
- `C:\Users\avidu\OneDrive\Documents\claude-sync\memory\khelsutra-guru\session-handoff.md`

I did **not** edit the live tracker inside `badminton-highlight-indexer` because you explicitly told me not to make changes in that repo.

****Leftover Items****

Item	Area	Status	Why still open
--- --- ---: ---			
`P13`	`sports-data-collector`	Gated	Needs owner decision on `.gitignore` artifact policy. Existing dirty `.gitignore` hunk preserved.
`P27`	non-git / GCS metadata backup	Gated	Non-git portion needs owner/input; engine doc side excluded by no-engine-edits.

| `P28` | sports-data / non-git cleanup | Gated | Blocked on owner/F: mount availability. |
| `P8-P10`, `P15`, `P23-P26` | engine / `badminton-highlight-indexer` | Excluded | Current
instruction says not to touch `badminton-highlight-indexer` or `main.py` refactoring. |

Everything else currently actionable in the allowed repos has been merged and reflected in the tracker memory.

573. user (2026-07-02T13:57:13.894Z)

yes now if you sync to head you can edit the files in the badminton highlight repo as well and I want you to complete all the pending items and treat the main.py refactoring as completed.

574. assistant (2026-07-02T13:57:30.972Z)

I'll resync the repos first and then re-read the live tracker plus the handoff/current-state docs so I'm working from the latest ground truth. Since you've reopened `badminton-highlight-indexer`, I'll include it again and treat the `main.py` refactor row as already complete in the tracker.

575. assistant (2026-07-02T13:57:44.313Z)

I found the handoff now says the old remediation-plan file was absorbed into the umbrella audit doc in the engine repo. I'm pulling each repo before reading anything else, then I'll use the umbrella tracker as the source of truth and reconcile it with the now-merged `main.py` work.

576. assistant (2026-07-02T13:58:00.259Z)

All pulls fast-forwarded cleanly. I'm reading the tracker rows and repo states now, with an eye toward separating ?already completed but not marked? from rows that still need an actual PR.

577. assistant (2026-07-02T13:58:11.468Z)

The umbrella tracker is still carrying some stale ?Todo? cells even though engine #431/#432 landed. I'm going to pull out the exact tracker table next, then start with the remaining low-risk engine rows in order.

578. assistant (2026-07-02T13:58:25.620Z)

The table itself is stale in several rows: P12 and P14?P22 are already merged elsewhere, and P23 needs to be closed out per your latest instruction. I'll still keep one small PR per real remaining work row, but I'll include a tracker reconciliation PR so the board stops lying to the next session.

579. assistant (2026-07-02T13:58:42.877Z)

For P8 I'm checking the existing upload API and job model before editing, because this one is about cleanup timing and I want the wrapper to be tiny and testable rather than a background service with hidden behavior.

580. assistant (2026-07-02T13:58:55.357Z)

Interesting: the cleanup tool already understands upload retention and has unit coverage; P8's missing piece is the packaged scheduler/wrapper plus docs. I'll add that as a small engine PR rather than rewriting the existing cleanup logic.

581. assistant (2026-07-02T13:59:12.234Z)

I'm going to cut the P8 branch from `origin/master` now and make only the upload-retention wrapper/docs/test changes on it. After that I'll run focused tests plus the repo gates before opening and merging the PR.

582. assistant (2026-07-02T13:59:24.848Z)

I'm adding a small `tools.upload_retention` entry point that reuses the existing cleanup planner, so the destructive decision logic remains in one place. The new code will dry-run by default and only delete when a scheduler passes `--apply`.

583. assistant (2026-07-02T14:00:48.558Z)

Focused tests pass; ruff only wanted import formatting in the new wrapper. I'm applying that mechanical cleanup, then I'll run the full engine verification set for the P8 branch.

584. assistant (2026-07-02T14:01:00.159Z)

The P8 branch is ready for gates. I'm running the full covered suite, ruff, mypy, and markdown link check now; this may take a bit, and I'll keep the exact coverage number for the PR body.

585. assistant (2026-07-02T14:03:31.185Z)

P8 gates are green: full suite `1586 passed, 11 skipped`, coverage `85.17%` against the `70%` floor, plus ruff/mypy/link-check clean. I'm re-verifying branch and HEAD now, then I'll stage only the P8 files and open the PR.

586. assistant (2026-07-02T14:03:34.863Z)

Branch and HEAD check out: this branch is still exactly on `origin/master` before my P8 commit. I'm staging the five intended P8 paths explicitly.

587. assistant (2026-07-02T14:03:47.600Z)

The P8 commit contains only the intended wrapper/docs/test files and `origin/master..HEAD` lists just my commit. I'm pushing it and opening the PR with the local verification recorded.

588. assistant (2026-07-02T14:04:17.751Z)

P8 is up as engine PR #433. While its checks run, I'm moving to the independent `sports-obsreport` P11 row so we can keep momentum without stacking engine work.

589. assistant (2026-07-02T14:07:13.379Z)

Engine PR #433 is fully green in GitHub Actions. I'm merging it now, then I'll fast-forward the local engine checkout before doing more engine rows.

590. assistant (2026-07-02T14:07:29.672Z)

P8 is merged as engine #433 and local master is fast-forwarded to `ab2e598`. I'm continuing on `sports-obsreport` P11 now: adding CI and fixing the customer-facing pluralization issue.

591. assistant (2026-07-02T14:07:40.022Z)

I'm reading the small report-rendering package now. The target looks narrow: `_plural("match", 3)` should be natural English, and the repo needs a basic workflow that runs lint and coverage.

592. assistant (2026-07-02T14:07:46.155Z)

I found the pluralizer is the simple `noun + "s"` rule, so `match` currently becomes `matches`. I'm branching from `origin/main`, then I'll add the irregular/`ch` plural rule, a regression test, and CI that runs ruff plus pytest coverage.

593. assistant (2026-07-02T14:08:05.742Z)

Baseline coverage is `96%`, so I'll set the new CI floor at `96%` and fix the existing ruff findings as part of this row. That keeps the repo from adopting a lint gate that starts red.

594. assistant (2026-07-02T14:08:30.413Z)

The P11 edit is in: CI workflow, ruff in dev extras, the plural regression, and cleanup for existing ruff findings. I'm running the local lint and coverage gate now.

595. assistant (2026-07-02T14:08:41.316Z)

P11 local verification is green: ruff passes and coverage is `96.48%` against the new `96%` floor. I'm checking branch/HEAD and the exact file set before staging.

596. assistant (2026-07-02T14:08:45.105Z)

The branch is still based directly on `origin/main` with no prior commits. I'm staging only the CI file, pluralizer/test, and the two ruff-cleanup test/import edits.

597. assistant (2026-07-02T14:08:59.553Z)

The P11 commit has only the intended six files and `origin/main..HEAD` is just this commit. I'm pushing and opening the PR now.

598. assistant (2026-07-02T14:09:31.067Z)

Sports-obsreport PR #1 is open. While its new CI runs, I'm moving back to the engine supply-chain rows; P9 is independent, and I'll wait for the obsreport PR to merge before pinning engine's obsreport dependency to the merged commit.

599. assistant (2026-07-02T14:09:59.631Z)

The tracker tag resolves to `1bf611d95262450c04fcf093334dd4c7116a2df7`. For P9 I'm going to centralize the git requirement in a tiny requirements file, then make CI, nightly, Docker, and the bootstrap script consume that instead of each carrying their own mutable tag literal.

600. assistant (2026-07-02T14:10:11.788Z)

Sports-obsreport #1 merged cleanly. I'm cutting P9 from the latest engine `origin/master` now and will keep it to the tracker dependency pin only.

601. assistant (2026-07-02T14:10:18.656Z)

I'm replacing the five duplicated `trackers` install literals with one root requirements file pinned to the resolved 40-character commit SHA, then pointing CI/nightly/Docker/bootstrap at that file.

602. assistant (2026-07-02T14:11:09.090Z)

That first patch missed the audit changelog context after the P8 merge, so no files changed. I'll apply it in smaller pieces.

603. assistant (2026-07-02T14:11:45.662Z)

The functional P9 files are patched. I'm updating the umbrella tracker row and changelog separately against the current merged audit doc.

604. assistant (2026-07-02T14:12:13.174Z)

The P9 tracker row is now marked done on the branch. I'm checking for any remaining functional mutable-tag install literals, then I'll run the engine gates.

605. assistant (2026-07-02T14:12:41.316Z)

The new requirements file resolves correctly in pip dry-run to the pinned SHA. I'm running the full engine gate again now; this dependency row still needs the same coverage discipline because it changes CI/deploy install paths.

606. assistant (2026-07-02T14:15:13.755Z)

Full engine gate is green again for P9: `1586 passed`, coverage `85.17%`, ruff/mypy/link-check clean. I'm doing one last literal scan with explicit filenames, then I'll commit.

607. assistant (2026-07-02T14:15:21.030Z)

The scan found no remaining functional `trackers@2.5.0` install literal. Branch and HEAD are cleanly based on `origin/master`; I'm staging only the P9 files now.

608. assistant (2026-07-02T14:15:32.364Z)

The P9 commit contains only the intended seven files and only my commit is ahead of `origin/master`. I'm pushing and opening the PR with the coverage and dependency dry-run noted.

609. assistant (2026-07-02T14:16:03.238Z)

Engine P9 is up as #434. While it runs CI, I'm checking the sports-data collector `.gitignore` hunk for P13; your latest `complete all pending?` gives enough direction to resolve it deliberately if the hunk is correct and manifest-preserving.

610. assistant (2026-07-02T14:19:03.996Z)

P9 is green in GitHub Actions too. I also inspected the collector `.gitignore` hunk: it's close, but `artifacts/` would hide `manifest.json` as well, so I'll turn it into an explicit `ignore` artifact payloads, unignore manifests? pattern for P13.

611. assistant (2026-07-02T14:19:18.360Z)

Engine P9 is merged as #434 and local master is at `14cdc2f`. I'm branching sports-data-collector from `origin/master` now and preserving the existing `.gitignore` hunk while tightening it.

612. assistant (2026-07-02T14:19:24.033Z)

P11 is also merged in sports-obsreport as #1; I'll reflect that in the central tracker on the next engine tracker-touching PR.

613. assistant (2026-07-02T14:19:31.967Z)

The pre-existing hunk was carried onto the P13 branch. I'm changing it from a broad `artifacts/` ignore to a payload-only ignore that lets bundle manifests remain visible to git.

614. assistant (2026-07-02T14:19:44.253Z)

The collector ignore policy is now manifest-preserving. I'm checking the ignore behavior explicitly, then I'll run the repo's lint and coverage gates.

615. assistant (2026-07-02T14:20:00.528Z)

`git check-ignore -v` shows the manifest path matching the negation rule, which is exactly what we want; my first boolean wrapper treated any match as failure. I'm using the verbose pattern output as the check here, then I'll run the actual collector CI commands.

616. assistant (2026-07-02T14:20:16.043Z)

P13 is verified: ruff clean, collector coverage `82.02%` against the `80%` floor, and the ignore rules hide weights while exposing `manifest.json`. I'm rechecking branch/HEAD, then staging only `.gitignore`.

617. assistant (2026-07-02T14:20:22.237Z)

This branch is still directly on `origin/master` with no commits yet. I'm committing the one intended `.gitignore` change.

618. assistant (2026-07-02T14:20:32.514Z)

The P13 commit contains just `.gitignore` and only my commit is ahead. I'm opening the sports-data-collector PR now.

619. user (2026-07-02T14:20:36.559Z)

(absolutely, you are now responsible to complete all the work related to this audit and present a final state after completing all the tasks ensuring all the docs are refreshed and uptodate)

620. assistant (2026-07-02T14:20:42.368Z)

Got it. I'll carry the audit through to a final reconciled state: every actionable row completed or explicitly marked blocked/owner-gated, all PRs merged green, and the umbrella tracker plus memory handoff/current-state refreshed at the end. Continuing with P13 PR creation now.

621. assistant (2026-07-02T14:21:02.133Z)

P13 is open as sports-data-collector #59. While its CI runs, I'm starting engine P10 now that sports-obsreport #1 is merged, so the engine can pin to the merged canonical commit instead of the old redirect/tag.

622. assistant (2026-07-02T14:21:34.913Z)

P13's CI is green. I'm starting engine P10 from fresh `origin/master`: update the engine's `reporting` extra and docs comments from the stale redirect/tag to the canonical `Khelsutra/sports-obsreport` commit that now includes PR #1.

623. assistant (2026-07-02T14:21:47.824Z)

I'm pinning to sports-obsreport commit `a0c7252c6e4159f8702360bea1cb46e5148b5c9f`, the merge commit for its CI/plural fix. That makes the optional extra immutable and removes the old `avidullu` redirect.

624. assistant (2026-07-02T14:23:33.928Z)

I'm picking up from the compacted state: P10 is mid-flight in the engine branch, and sports-data #59 is ready to merge once I re-check it. I'll finish those first, then resync the tracker against remote head before moving into the remaining audit rows.

625. assistant (2026-07-02T14:23:52.587Z)

The sports-data PR is clean and its CI passed, so I'm merging that while the engine P10 branch gets its local verification. The engine branch is still exactly on top of the last merged master commit, with only the intended P10 files modified.

626. assistant (2026-07-02T14:26:09.187Z)

Good: sports-data merge returned cleanly. The obsreport dry run hit local SSH host-key verification before dependency resolution, so I'm going to re-run that with a one-command SSH host-key acceptance setting and give the engine suite a longer timeout; the suite was still running when the first ceiling cut it off.

627. assistant (2026-07-02T14:28:36.358Z)

The full engine suite is green at the same 85.17% coverage floor result as the previous engine waves. One useful catch: this local Windows environment cannot clone the private obsreport repo over SSH because there's no GitHub SSH key, so I'm checking whether HTTPS can resolve the canonical pinned repo cleanly before I decide whether to adjust the dependency URL.

628. assistant (2026-07-02T14:28:50.066Z)

sports-data is now fast-forwarded to the merged artifact-manifest fix. The engine static checks

are clean; I'm doing one dependency-resolution check over HTTPS now because that would make the canonical obsreport pin easier to verify locally and in CI than the old SSH form.

629. assistant (2026-07-02T14:29:08.030Z)

HTTPS resolves the pinned obsreport commit cleanly and avoids the local SSH-key trap, so I'm going to switch the P10 dependency from canonical SSH to canonical HTTPS while keeping the immutable commit pin. That keeps the fix smaller and more reproducible.

630. assistant (2026-07-02T14:29:11.544Z)

I've got the exact P10 diff in view. Next edit is surgical: update only the obsreport URL/comment wording to the HTTPS direct reference that just resolved in pip, then re-run the small verification set that can be affected by that change.

631. assistant (2026-07-02T14:29:28.701Z)

One stale backlog sentence still says "pin git+ssh tag later," which is now false after P10/P12. I'm cleaning that while I'm here so the docs don't carry a completed task as future work.

632. assistant (2026-07-02T14:32:03.474Z)

P10 is verified now: dependency dry-run resolves over HTTPS to the exact commit, full engine suite is green at 85.17%, and ruff/mypy/docs links are clean. I'm doing the required branch/HEAD re-check before staging the explicit files.

633. assistant (2026-07-02T14:32:07.233Z)

The required pre-stage check is clean: branch is `fix/pin-obsreport-dependency`, HEAD is still the P9 merge commit, and there are no local commits riding above origin/master yet. Staging only the four P10 files now.

634. assistant (2026-07-02T14:32:14.405Z)

P10 is committed as a single local commit. I'm checking the branch delta against `origin/master`, then I'll push/open the PR with the verification record in the body.

635. assistant (2026-07-02T14:32:18.486Z)

The branch delta is exactly one P10 commit and four intended files. Pushing it now and opening the PR.

636. assistant (2026-07-02T14:32:24.207Z)

Push is up. I'm opening the PR with the local coverage number and the dependency dry-run called out, since that was the meaningful verification for this wave.

637. assistant (2026-07-02T14:32:36.342Z)

P10 is PR #435. I'm watching the GitHub checks before merging, then I'll fast-forward local master and move into the remaining rows.

638. assistant (2026-07-02T14:32:41.947Z)

GitHub hasn't attached checks to #435 yet, so I'm going to poll the PR metadata rather than assume. If the workflows stay absent, I'll treat that as a GitHub dispatch delay and wait a bit before merging.

639. assistant (2026-07-02T14:33:02.027Z)

Checks are attached now; the quick lanes are green and the full/golden/license/build lanes are still running. I'm waiting for the whole CI rollup before merging.

640. assistant (2026-07-02T14:35:46.996Z)

All #435 checks are green. I'm merging it and fast-forwarding the local engine master before starting the next item.

641. assistant (2026-07-02T14:36:07.970Z)

P10 is merged as #435 and local engine master is fast-forwarded to `ea9f0b2`. I'm going to resync the workspace/memory state now and rebuild the pending list from the live tracker plus current repo heads, because the remaining work depends on what other sessions have already landed.

642. assistant (2026-07-02T14:36:19.216Z)

The root is a non-git workspace, so there was nothing to pull there. The handoff/current-state are stale relative to the PRs just merged, so I'll refresh them later; first I'm pulling each sub-repo head before reading their current audit/tracker state.

643. assistant (2026-07-02T14:36:35.287Z)

All five repos are synced to remote head. I'm reading the live tracker rows now and cross-checking them against the actual merged PRs so the next fixes don't duplicate already-closed work.

644. assistant (2026-07-02T14:36:45.736Z)

The live tracker confirms the status drift: P11-P22 are mostly already merged but still marked Todo in the engine doc. The real remaining work appears to be P15, P25, P26, P27, and possibly P28 if the F: source drive is mounted; P23/P24 are to be treated as complete per your latest instruction.

645. assistant (2026-07-02T14:36:53.418Z)

I'm starting P15 as its own engine PR: branch from `origin/master`, inspect the existing UI bundle workflow, then rebuild/add a durable staleness signal instead of hand-copying files.

646. assistant (2026-07-02T14:37:04.979Z)

P15 branch is cut from `origin/master` at `ea9f0b2`. I'm reading the bundle/build conventions now, including any existing UI metadata so the refresh is repeatable.

647. assistant (2026-07-02T14:37:18.844Z)

The existing `scripts/bundle\_ui.sh` already captures repo commit, engine API version, content hash, and notices. I'm using that path to rebuild `backend/ui` from khelsutra `61e5ccb`, then I'll add a CI-readable staleness check if one doesn't already exist.

648. assistant (2026-07-02T14:37:39.519Z)

The vendored bundle is rebuilt from khelsutra `61e5ccb`. I'm adding a lightweight check that reports whether `backend/ui/ui\_version.json` matches a sibling khelsutra checkout when one is present, so future drift is visible without requiring CI to clone another private repo.

649. assistant (2026-07-02T14:37:47.948Z)

There are already tests for `?bundle exists?` and API_VERSION skew. I'm extending that area with a side-effect-free helper and unit tests so the doctor can warn when a sibling khelsutra checkout has moved ahead.

650. assistant (2026-07-02T14:37:59.358Z)

I'm making the doctor warning testable without shelling out in the tests: the check reads `ui\_version.json`, optionally accepts a known source commit for tests, and in normal use

compares against `KHELSUTRA\_DIR` or `../kheksutra`.

651. assistant (2026-07-02T14:38:28.566Z)

The doctor check is in place with tests. I noticed the import ordering may trip ruff, so I'm tidying that now, then I'll update the tracker/changelog for P15 before running the full gate.

652. assistant (2026-07-02T14:38:33.917Z)

I'm recording P15 in the live tracker on this branch, but I'll leave the broader P11-P24 reconciliation for the final tracker PR so each change stays reviewable.

653. assistant (2026-07-02T14:39:03.167Z)

P15 now has the code and tracker update. I'm running the targeted Python tests first, then the full engine coverage gate and static checks; after that I'll do a quick visual validation of the engine-served bundle into the existing screencast evidence folder.

654. assistant (2026-07-02T14:41:41.695Z)

The Python gates are green with coverage unchanged at 85.17%. For the UI evidence, I'm launching the engine with a temporary app home, loading the bundled SPA through the actual FastAPI server, and saving a screenshot under the existing `kheksutra-screencast/\_raw` folder.

655. assistant (2026-07-02T14:42:05.017Z)

The bundled SPA rendered through the engine and the screenshot landed at `kheksutra-screencast/\_raw/p15-bundled-ui-20260702-201156.png`. I'm doing one visual sanity check of the screenshot file before finalizing the PR.

656. assistant (2026-07-02T14:42:19.453Z)

The screenshot looks good: the engine-served bundle renders the current Studio UI, not a blank/static shell. I'm doing the last whitespace/status checks and the real sibling-check invocation before committing.

657. assistant (2026-07-02T14:42:31.279Z)

P15's actual bundle check passes against the sibling kheksutra checkout, and the new bundle contains the current library UI code. I'm doing the branch/HEAD safety check, then staging the explicit bundle/script/test/tracker paths including the two old hashed assets as explicit removals.

658. assistant (2026-07-02T14:42:34.130Z)

The branch is cleanly based on `origin/master` with no local commits yet. Staging explicitly now.

659. assistant (2026-07-02T14:42:39.779Z)

Only the intended P15 paths are staged. I'm committing this as the vendored SPA refresh.

660. assistant (2026-07-02T14:42:48.314Z)

P15 is committed. I'm confirming the branch contains only that commit over master, then I'll push and open the PR with test and screenshot evidence.

661. assistant (2026-07-02T14:42:51.168Z)

The branch delta is a single P15 commit. Pushing and opening the PR now.

662. assistant (2026-07-02T14:42:57.467Z)

Push is up. I'm including the screenshot path and coverage result in the PR description so the UI refresh has concrete evidence attached.

663. assistant (2026-07-02T14:43:11.959Z)

P15 is PR #436. I'm watching CI, then I'll merge and move on.

664. assistant (2026-07-02T14:43:17.181Z)

GitHub hasn't attached checks to #436 immediately, same as the previous PR. I'm polling the PR rollup after a short delay.

665. assistant (2026-07-02T14:43:38.196Z)

#436 checks are running; quick lint/leak/docs lanes already passed. I'm waiting on the full/golden/build lanes.

666. assistant (2026-07-02T14:46:24.482Z)

All #436 checks are green. I'm merging P15 and fast-forwarding local master.

667. assistant (2026-07-02T14:46:35.154Z)

P15 is merged as #436 and local engine master is at `d490485`. Next I'm going after P25, but I'll first inspect the database module shape carefully so the split stays behind the existing facade and doesn't churn call sites.

668. assistant (2026-07-02T14:46:41.233Z)

For P25, I'll keep the facade promise: `backend.infrastructure.database.Database` remains the public class, and any split should be internal/mixin-style or repository-style with zero call-site churn. Branching from `origin/master` now.

669. assistant (2026-07-02T14:46:57.279Z)

P25 branch is cut at the latest master. I'm reading `database.py` and its tests now to find the lowest-risk split: domain methods out, schema/migration core staying put.

670. assistant (2026-07-02T14:47:01.643Z)

The first read confirms the Database class mixes schema core, video/segment/candidate CRUD, jobs/cost/policy, and user-model store. I'm using AST line ranges now so I can move whole methods mechanically rather than rewriting SQL.

671. assistant (2026-07-02T14:47:29.238Z)

The split will be mixin-style: new domain modules provide media/segment, job/cost/policy, and user-model methods; `Database` inherits them and keeps schema/migration/connection logic. That preserves `Database(...)` as the facade and avoids changing any caller.

672. assistant (2026-07-02T14:47:47.692Z)

I'm going to do this as a mechanical method move: exact method blocks leave `database.py`, new mixin modules get those same methods plus small core-method stubs for type checking. Afterward I'll run focused DB tests, mypy, and the full coverage gate.

673. assistant (2026-07-02T14:47:56.407Z)

I'm making the mechanical split now: media CRUD, jobs/cost/policy, and user-model store move to their own modules; the schema ladder and connection helpers remain in `database.py`.

674. assistant (2026-07-02T14:48:11.073Z)

The method blocks are now moved, but `Database` still needs to inherit the mixins and drop imports it no longer owns. I'm applying that small hand edit next.

675. assistant (2026-07-02T14:48:24.777Z)

The public class now inherits the domain mixins. I'm adding a tiny facade test so future edits don't accidentally make callers import the split modules directly.

676. assistant (2026-07-02T14:48:40.271Z)

The facade test is in. I'm running focused database tests and static checks now; this is the spot where any mechanical-move mistake should show quickly.

677. assistant (2026-07-02T14:48:49.480Z)

The focused DB behavior tests all passed; ruff only found a now-unused `json` import left in the schema core. I'm removing that and then updating the code map/tracker for the split.

678. assistant (2026-07-02T14:49:05.599Z)

CODE_MAP still describes the old monolith and an outdated schema version, so I'm correcting that as part of P25.

679. assistant (2026-07-02T14:49:30.212Z)

I'm marking P25 done in the live tracker and noting the database split in the audit changelog.

680. assistant (2026-07-02T14:49:49.983Z)

The database split is now documented and tracker-marked. I'm running the full engine gate next: full coverage suite, ruff, mypy, markdown links, and diff check.

681. assistant (2026-07-02T14:52:19.758Z)

The gate passed, but coverage dipped from the previous 85.17% to 85.16% because the new mixin core-method stubs count as uncovered lines. That is exactly the sort of tiny regression you asked me to avoid, so I'm excluding those non-runtime stubs from coverage and re-running the coverage gate.

682. assistant (2026-07-02T14:52:36.920Z)

The stubs are now marked as facade-implemented and excluded from coverage. Re-running the full verification set so the PR body can honestly report the final coverage.

683. assistant (2026-07-02T14:55:04.809Z)

Coverage is back above the pre-split number at 85.20%, so no regression. I'm also running an AST equality check for the moved database methods against `origin/master` to prove this was a mechanical split, not a SQL rewrite.

684. assistant (2026-07-02T14:55:17.721Z)

The AST check confirms 32 moved database methods are identical to `origin/master`. I'm doing the required branch/HEAD check, then staging only the P25 files.

685. assistant (2026-07-02T14:55:21.319Z)

Branch safety is good: no commits on top of origin/master yet, and only P25 files are modified/untracked. Staging explicit paths now.

686. assistant (2026-07-02T14:55:26.549Z)

Only explicit P25 paths are staged. Committing the database facade split.

687. assistant (2026-07-02T14:55:33.785Z)

P25 is committed. I'm confirming the branch delta, then I'll push and open the PR with AST-identical proof and coverage.

688. assistant (2026-07-02T14:55:36.765Z)

The branch has exactly one P25 commit. Pushing and opening the PR.

689. assistant (2026-07-02T14:55:43.640Z)

Push is up. The PR body will call out both behavior tests and the AST equality check so this reads as a controlled refactor.

690. assistant (2026-07-02T14:55:55.245Z)

P25 is PR #437. I'm waiting for GitHub checks before merging.

691. assistant (2026-07-02T14:58:59.266Z)

All #437 checks are green. I'm merging P25 and fast-forwarding local master, then I'll move to P26 issue cleanup.

692. assistant (2026-07-02T14:59:13.382Z)

P25 is merged as #437 at `66e8ceb`. For P26 I'm collecting the referenced issues across the engine and SPA repos, then I'll close only the ones that are clearly superseded by merged PRs and comment where residual scope remains.

693. assistant (2026-07-02T14:59:28.595Z)

I'm querying both likely issue owners now; the numbers overlap across repos, so I'm not assuming the engine repo owns every referenced issue.

694. assistant (2026-07-02T15:00:05.000Z)

I found the stale issue set: several khelsutra UX issues are now fully covered by merged PRs #144-#148; a few engine issues are partly code-done but still have owner/live-infra residuals. I'm going to close only completed/superseded issues and leave narrowed comments on the remaining ones.

695. assistant (2026-07-02T15:00:30.754Z)

The completed khelsutra UX duplicates are now closed with PR references. I'm adding scoped comments to the issues that are still real but narrowed, and closing the remaining main.py refactor follow-up per your explicit 'treat completed / don't work on main.py' directive.

696. assistant (2026-07-02T15:00:58.426Z)

P26 issue actions are done: completed duplicates closed, owner/live-infra residuals left open with narrowed comments, and #388 closed as not planned per your directive. I'm verifying the remaining open issue set before marking P26 in the tracker later.

697. assistant (2026-07-02T15:01:03.997Z)

P26 is functionally complete. Remaining open issues are genuine residual product/owner-live-infra items, not stale duplicates; I'll reflect that in the final tracker reconciliation. Next is P27: backing up the Roora non-video artifacts and updating `DATA\_IN\_GCS.md`.

698. assistant (2026-07-02T15:01:13.401Z)

For P27 I'll branch a small engine docs PR, but the actual payload is a GCS archive copy of non-video Roora artifacts. I'm first listing the local artifact set and current bucket archive so I

can exclude videos and avoid duplicate/broad uploads.

699. assistant (2026-07-02T15:01:28.506Z)

P27 branch is cut from `origin/master` at `66e8ceb`. Now I'm enumerating the Roora artifact folders and the bucket prefix before copying.

700. assistant (2026-07-02T15:02:55.189Z)

I'll pick this up from the handoff and treat it as the active audit closeout now. I'm going to resync the root/repo state first, then work through the remaining rows with small PRs and keep the tracker honest as each one lands.

701. assistant (2026-07-02T15:02:59.149Z)

I'm using the GitHub publishing workflow skill for the PR/merge loop, since this task is explicitly about landing branches through GitHub while keeping branch hygiene tight.

702. assistant (2026-07-02T15:03:16.934Z)

I'm going to pull each repo I may touch before inspecting files. Since the engine repo is mid-task on a local branch from the prior session, I'll handle it carefully: fast-forward the branch from its remote base only if it is still clean and based on `origin/master`.

703. assistant (2026-07-02T15:03:23.455Z)

The sibling repos are current. I'm checking the engine checkout now, because it's the one with the active P27 branch and I don't want to accidentally drag a sibling branch into the next PR.

704. assistant (2026-07-02T15:03:37.049Z)

All repos I expect to touch fast-forward cleanly, so the context should now reflect remote head. I'm reloading the handoff and live tracker before making the next branch changes.

705. assistant (2026-07-02T15:03:47.019Z)

The old standalone remediation plan is gone on current engine master; the tracker was folded into the umbrella audit doc. I'm switching my reads to that file and the data map, which matches the newer repo state.

706. assistant (2026-07-02T15:03:54.817Z)

The tracker is stale in several rows even though the merged PRs are real: P11?P22, P26, and the ?this PR? placeholders need reconciliation. I'll fix P27 first, then make one tracker cleanup pass so it reflects landed state rather than half-day-old state.

707. assistant (2026-07-02T15:04:00.164Z)

For P27 I'm inventorying only the non-video Roora pilot artifacts and the small CSV/trajectory side files. The raw videos stay out of scope unless the owner explicitly approves a large upload.

708. assistant (2026-07-02T15:04:16.079Z)

The P27 payload is small: 730 non-video pilot files (~5.6 MB), two root CSVs, and seven Roora trajectory CSVs. I'll stage that exact set with a manifest, verify no video extensions slipped in, then upload it under the planned archive prefix.

709. assistant (2026-07-02T15:04:25.813Z)

Interesting twist: the Roora archive prefix already exists in GCS and appears to contain the pilot XML/annotation payload. I'm going to compare it against the local inventory, then only add missing low-risk sidecars/manifest instead of blindly duplicating objects.

710. assistant (2026-07-02T15:04:40.853Z)

The existing GCS prefix has 717 non-video objects and zero video objects, so the big-risk part is already handled. There are 22 local sidecar/metadata files not represented there, so I'm preparing a compact `metadata/` add-on plus an archive manifest rather than altering the existing 717-object payload.

711. assistant (2026-07-02T15:04:44.164Z)

I'm doing a path-level diff now. If the missing pilot files are only packaging metadata or feedback docs, I'll add them under a clear sidecar layout and document the exact counts.

712. assistant (2026-07-02T15:05:03.306Z)

The prior upload flattened one folder level for the pilot payload. That still preserved the annotation content, but the docs need to be honest about the chosen GCS layout; I'll add the missing feedback/spec/zip sidecars, root CSVs, Roora trajectories, and a manifest that records both local source paths and GCS objects.

713. assistant (2026-07-02T15:05:25.815Z)

I'm going to build the manifest from SHA-256 hashes of the local small files. This gives the archive an audit trail without moving the 4.8 GB raw videos.

714. assistant (2026-07-02T15:06:00.760Z)

The manifest says 739 small files in scope, with 717 already present and 22 missing. I'm uploading only those 22 plus `manifest.json`, then I'll verify the object count and that no video objects exist under the archive prefix.

715. assistant (2026-07-02T15:06:43.157Z)

P27's bucket write is complete: 23 staged files total, consisting of 22 missing non-video files plus the archive manifest. I'm verifying the prefix now before editing the docs.

716. assistant (2026-07-02T15:07:04.430Z)

I'm editing the engine docs now: `DATA\_IN\_GCS.md` gets the Roora archive map, and the audit tracker gets P27 marked done with the exact bucket prefix, object count, and ?no video objects? verification.

717. assistant (2026-07-02T15:07:15.435Z)

I'm keeping this PR scoped to P27: the data map plus the P27 tracker/changelog. The larger tracker reconciliation pass will come after P28 so the final doc state is not pretending ahead of the work.

718. assistant (2026-07-02T15:07:55.301Z)

The P27 doc changes are in place. I'm also updating the doc-status row so a future session sees that the umbrella tracker has progressed beyond the earlier Wave 2 snapshot.

719. assistant (2026-07-02T15:08:10.929Z)

Now I'm running the docs checks for the P27 branch. Since this is docs plus an external bucket archive, there's no code coverage surface here, but link/diff hygiene still matters.

720. assistant (2026-07-02T15:08:18.954Z)

Docs checks are clean. I'm re-verifying branch and HEAD before staging, then I'll stage only the three P27 doc paths.

721. assistant (2026-07-02T15:08:22.940Z)

The branch is cleanly on `data/archive-roora-pilot-artifacts` at engine master head `66e8ceb`, with no extra commits above `origin/master`. I'm committing the P27 doc update now.

722. assistant (2026-07-02T15:08:31.602Z)

P27 is committed. I'm doing the pre-PR ancestry check next, then I'll push and open the PR with the bucket verification included in the body.

723. assistant (2026-07-02T15:08:43.793Z)

The P27 branch is pushed. I'm opening the PR as a normal PR so it can be merged once checks are green, with the docs-only/coverage-not-applicable note spelled out.

724. assistant (2026-07-02T15:08:58.070Z)

PR #438 is open. I'm watching the checks now; after they're green I'll merge it and resync master before moving on.

725. assistant (2026-07-02T15:11:52.232Z)

All checks on #438 are green, including the full offline suite and markdown link check. I'm merging it now, then I'll pull master back to the local checkout before starting P28.

726. assistant (2026-07-02T15:12:09.507Z)

P27 is merged as engine master `0f32d39`. I'm cleaning up the local branch state and pulling master before the P28 check.

727. assistant (2026-07-02T15:12:15.759Z)

P27 is done and merged. Next I'm checking whether the `F:` source drive is actually mounted and where the six owned golden source videos live before deciding whether P28 can be completed or must stay explicitly blocked.

728. assistant (2026-07-02T15:12:25.744Z)

I'll pull sports-data again before reading it for P28, since that repo may be where the durable metadata belongs. Then I'll inspect the existing corpus/metadata conventions instead of inventing a random sidecar.

729. assistant (2026-07-02T15:12:38.430Z)

Good news: `F:` is mounted. P28 can be completed, not just marked blocked. I'm locating the existing collector metadata/data conventions now so the hashes land somewhere durable and reviewable.

730. assistant (2026-07-02T15:12:55.713Z)

The audit's P28 points at sports-data-collector's registry: those six rows have stale `local\_path`s and no durable hashes. I'm inspecting the SQLite schema now; if there's no hash column, I'll add a small committed JSON manifest beside the DB rather than mutating schema just for a one-off audit record.

731. assistant (2026-07-02T15:13:08.679Z)

The nominal F: path from the script is stale on this machine, so I'm searching the mounted drive for the six exact MP4 filenames. That's still P28's job: verify the actual source paths before hashing.

732. assistant (2026-07-02T15:13:22.759Z)

I found all six MP4s under `F:\[Khelsutra] GoPro Hero Black 12\GoPro Recorded Videos\Training\Golden Labelled`, while the DB still points at the old deleted

`C:\Users\avidu\Projects\Annotation Setup\Collect\...\` path. I'm going to hash the six exact files and then update the collector registry plus a small committed hash manifest.

733. assistant (2026-07-02T15:13:25.453Z)

Hashing ~39 GB will take a little while. I'm using MD5 because that's the downstream indexer join key and the P28 acceptance criterion; I'll avoid a second SHA-256 pass unless the repo already expects it.

734. assistant (2026-07-02T15:15:19.124Z)

The hashes are captured. I'm cutting the P28 collector branch from `origin/master` now, then I'll update the DB rows, add `data/golden\_source\_videos.json`, and add a lightweight test that keeps the JSON and DB paths in sync.

735. assistant (2026-07-02T15:15:27.589Z)

I'm updating the SQLite registry through SQLite itself, not byte-editing the DB. The new committed JSON is the durable hash record; the DB path fix makes future exports hashable when `F:` is mounted.

736. assistant (2026-07-02T15:15:39.271Z)

The six DB `local\_path` values now point at the mounted F: source files. I'm adding the hash manifest and a regression test to keep that file and the DB aligned.

737. assistant (2026-07-02T15:16:14.015Z)

I'm checking the actual DB string values before running tests; Windows path escaping is one of those tiny details that can bite when a binary registry meets JSON.

738. assistant (2026-07-02T15:16:24.436Z)

The DB paths are correctly stored as normal Windows paths. I'm running the focused test plus the collector's lint and coverage gate now.

739. assistant (2026-07-02T15:16:32.585Z)

The first full coverage invocation used a stale global `pytest` entry point that resolved tests from an older checkout and couldn't import `src`. I'm rerunning through the current Python module entry point from this repo, which is the command that avoids that path skew.

740. assistant (2026-07-02T15:16:57.270Z)

The current-repo coverage run passed: 290 passed, 2 skipped, total coverage 82.02% against the 80% floor. I'm checking the exact diff now, including that the old collector `.gitignore` is not riding along.

741. assistant (2026-07-02T15:17:07.592Z)

I'm adding one README pointer too, so the hash manifest is discoverable from the collector docs instead of being a quiet data file.

742. assistant (2026-07-02T15:17:21.897Z)

The README now points at the committed hash manifest. I'm rerunning the current-repo checks after that doc touch, then I'll stage explicit paths.

743. assistant (2026-07-02T15:17:37.808Z)

Checks are green and coverage is still 82.02% over the 80% floor. I'm doing the branch/HEAD recheck and then staging only `README.md`, the DB, the JSON manifest, and the new test.

744. assistant (2026-07-02T15:17:41.304Z)

The pre-stage check is clean: branch is `data/golden-source-video-hashes`, HEAD is remote master, and no sibling work is present. I'm staging the four intended paths only.

745. assistant (2026-07-02T15:17:49.370Z)

P28 is committed locally. I'm checking the branch ancestry and then I'll open the collector PR with the hash list and coverage result in the body.

746. assistant (2026-07-02T15:17:58.849Z)

The P28 branch is pushed. I'm creating the PR in sports-data-collector now.

747. assistant (2026-07-02T15:18:16.473Z)

Sports-data-collector PR #60 is open. I'm watching its CI and will merge it once green.

748. assistant (2026-07-02T15:18:44.059Z)

PR #60 CI is green. I'm merging it, then pulling sports-data master back to head.

749. assistant (2026-07-02T15:18:54.556Z)

P28 is merged as sports-data-collector master `8303fd6`. I'm returning that checkout to master and then I'll update the engine tracker to show P27/P28 and all earlier rows accurately.

750. assistant (2026-07-02T15:19:02.050Z)

The last two open work rows are now merged. I'm moving into the final tracker reconciliation: engine audit doc, data map, doc status, and memory files all need to match the merged reality.

751. assistant (2026-07-02T15:19:06.850Z)

Before touching the engine tracker again, I'm pulling engine master so the final docs branch starts from #438.

752. assistant (2026-07-02T15:19:39.305Z)

The final tracker pass will be a docs-only engine PR. I'm going to mark rows by actual merged PR numbers, retire the main.py/P24 scope per your instruction, add P28's sports-data PR, and leave only owner-gated/residual work explicitly listed.

753. assistant (2026-07-02T15:20:05.059Z)

One more bit surfaced during P28: the Phase-2 copy helper's `F:` root still points at the old folder name. I'm going to fix that alongside the final tracker reconciliation so the docs do not point future sessions at a stale helper.

754. assistant (2026-07-02T15:20:14.276Z)

That copy helper exists locally but is not listed by `git ls-files`, so I'll leave it alone rather than add scratch tooling in the final tracker PR. I'll make the durable docs point at the actual verified `F:` path and the sports-data hash manifest instead.

755. assistant (2026-07-02T15:20:23.769Z)

I'm editing the final tracker now. The point is to make it boringly factual: actual PR numbers, actual merged states, and only the remaining owner-gated choices left in the residual list.

756. assistant (2026-07-02T15:22:33.095Z)

The umbrella tracker is reconciled. I'm updating the data map and doc indexes to point at the

P28 hash manifest and reflect the audit as done-with-residuals, rather than in progress.

757. assistant (2026-07-02T15:22:52.894Z)

The data map now records the P28 hash manifest and the verified F: root. I'm updating the docs index/register next so the "start here" trail says the board is closed, not still mid-wave.

758. assistant (2026-07-02T15:23:16.630Z)

The index/register now says "done with residuals." I'm scanning for stale placeholders like "this PR", "Todo", and "Blocked" before running the doc checks.

759. assistant (2026-07-02T15:24:29.097Z)

I'm picking back up at the tracker reconciliation branch. I'll verify the docs-only changes locally, then open and merge the final tracker PR before refreshing the shared handoff/current-state files.

760. assistant (2026-07-02T15:24:44.738Z)

The branch is exactly where expected: "docs/final-audit-tracker" on engine "0f32d39" from "origin/master", with only the four tracker/docs files modified. I'm running the docs checks now before committing.

761. assistant (2026-07-02T15:24:48.046Z)

The link checker is clean. "git diff --check" caught one trailing-space line in the audit doc, so I'm removing that little snag and then I'll re-run the checks.

762. assistant (2026-07-02T15:25:03.839Z)

That whitespace issue is fixed. I'm doing the commit gate now: link check, diff whitespace check, and the branch/HEAD verification right before staging.

763. assistant (2026-07-02T15:25:07.119Z)

The pre-stage gate is clean: no broken docs links and no diff whitespace errors. I'm staging only the four tracker/doc files now, then committing the final reconciliation.

764. assistant (2026-07-02T15:25:11.448Z)

The index contains only those four files. I'm committing this as a docs-only tracker finalization.

765. assistant (2026-07-02T15:25:17.631Z)

Commit "87598c9" is made. Before opening the PR, I'm checking that this branch contains only that one commit above "origin/master", then I'll push it.

766. assistant (2026-07-02T15:25:21.888Z)

The branch delta is clean: one commit, no stray working-tree changes. Pushing and creating the PR now.

767. assistant (2026-07-02T15:25:28.168Z)

The final tracker branch is on GitHub. I'm opening the PR with the verification notes, then I'll let CI tell us whether the docs-only reconciliation is clean.

768. assistant (2026-07-02T15:25:38.530Z)

PR #439 is open. Since the number is now known, I'm replacing the last "this PR" changelog marker with "#439" so the tracker doesn't keep a self-reference after merge.

769. assistant (2026-07-02T15:25:47.516Z)

That final self-reference is updated. I'm re-running the docs checks and the branch/HEAD pre-stage check before amending the PR commit.

770. assistant (2026-07-02T15:25:51.355Z)

Checks are clean again, and the branch has only the tracker commit above `origin/master`. I'm staging just the audit doc and amending the PR commit.

771. assistant (2026-07-02T15:25:56.022Z)

The amended PR branch still has a single commit, now `5f3e5f6`. I'm updating the remote branch with a force so CI runs against the exact final docs.

772. assistant (2026-07-02T15:26:02.475Z)

PR #439 is updated. I'm watching the GitHub checks now and will merge it once they go green.

773. assistant (2026-07-02T15:26:09.921Z)

GitHub hasn't attached checks to the branch yet, which sometimes happens for a brand-new docs-only PR before Actions wakes up. I'm checking the PR metadata and workflow state rather than assuming.

774. assistant (2026-07-02T15:26:14.043Z)

Actions are running now; the empty check rollup was just early. I'll watch the latest CI run and then merge #439 if it passes.

775. assistant (2026-07-02T15:29:29.193Z)

CI is green, including the full offline test suite with the coverage floor, ruff, mypy, wheel smoke, link check, and cross-OS golden fixtures. I'm merging #439 now.

776. assistant (2026-07-02T15:29:45.556Z)

#439 is merged to engine master as `c30ff22`. Next I'm refreshing the local checkout and then I'll update the shared memory handoff/current-state files so a new session lands on the real final state.

777. assistant (2026-07-02T15:29:56.476Z)

The primary repos are clean and on their expected default branches after the merge. I'm reading the shared handoff/current-state one more time before editing, so the memory update is a true refresh rather than a blind overwrite.

778. assistant (2026-07-02T15:30:36.965Z)

The engine tracker PR is merged and all repos are clean at their expected heads. I found the shared memory files still had earlier same-day "next step" text, so I'm going to rewrite them down to the final, honest handoff instead of letting stale resume cues linger.

779. assistant (2026-07-02T15:30:49.406Z)

There's also a stale root-level audit snapshot outside git. I'm replacing it with a short pointer to the live umbrella tracker in the engine repo, so the workspace root no longer advertises the pre-remediation audit as current.

780. assistant (2026-07-02T15:31:52.209Z)

The handoff/current-state are now short and final-state-oriented, and the root audit copy is no longer stale. I'm doing one last sanity pass across repo status, PR state, and the tracker rows

before I wrap this up.

781. assistant (2026-07-02T15:32:00.267Z)

The only remaining `this PR` hits in the live audit doc are in the original audit snapshot text and in the changelog sentence explaining that placeholders were replaced, not in live tracker rows. I'm checking repo cleanliness and open PR state now.

782. assistant (2026-07-02T15:32:20.835Z)

One more housekeeping item surfaced: a few earlier audit-era PRs are still open in `kheksutra` and `sports-data-collector` even though the later rows landed. I'm checking whether they're superseded or still carrying unique work before I close or act on them.

783. assistant (2026-07-02T15:32:41.385Z)

These are not random product PRs: #141 and sports-data #54-#56 are earlier audit PRs. I'm comparing their intended fixes against current master so I don't merge stale/conflicting work or close something with unique value.

784. assistant (2026-07-02T15:32:54.685Z)

Current master already contains the collector artifact policy, so sports-data #54 is likely superseded. The other open PRs still have some unique content, but a couple mention stale paths/prefixes, so I'm going to inspect diffs and mergeability rather than blindly merging old work.

785. assistant (2026-07-02T15:33:09.942Z)

The old sports-data #55 has a useful `curate verify` idea but also stale binary DB edits from before the F: hash pass. I'm checking the audit tracker language now to decide whether that old PR is audit-required or a future nice-to-have.

786. assistant (2026-07-02T15:33:25.004Z)

#55 is now conflicting because its DB edits predate the correct F: MD5/path fix in #60, so merging it would be unsafe. #56 is the remaining question: it may still add unique small golden-label fixtures, so I'm checking whether it can be refreshed or whether #438's GCS archive fully supersedes it.

787. assistant (2026-07-02T15:33:57.858Z)

#56 is clean/mergeable and adds unique small Roora golden-label fixtures, so I'm going to verify it on top of current sports-data master and merge it if the repo gates stay green. The conflicting #54/#55 will get closed as superseded rather than left dangling.

788. assistant (2026-07-02T15:34:02.562Z)

I'm testing #56 as a cherry-pick on current `origin/master`, which is the exact merge-result shape I care about. That avoids trusting an old branch head just because it happens to be mergeable.

789. assistant (2026-07-02T15:34:15.466Z)

The cherry-pick applies cleanly on current sports-data master. I'm running the focused data test, the full coverage gate, ruff, and diff whitespace check before merging #56.

790. assistant (2026-07-02T15:34:29.986Z)

#56's merge-result test is green: focused test passed, full suite is `293 passed, 2 skipped`, coverage still `82.02%` against the 80 floor, and ruff/diff checks are clean. I'm merging it now.

791. assistant (2026-07-02T15:34:40.995Z)

The merge command hit a cleanup wrinkle: GitHub couldn't delete a local branch because an old temp worktree still references it. I'm checking whether the PR itself merged; if it did, I'll just clean up local state without touching that worktree.

792. assistant (2026-07-02T15:34:46.743Z)

#56 is merged as `1423e84`; only local branch deletion failed because a stale temp worktree owns the old branch name. I'm moving the main checkout back to `master` and fast-forwarding it.

793. assistant (2026-07-02T15:35:01.980Z)

Sports-data #56 is now on master. I'm closing the two older conflicting sports-data PRs with explicit supersession notes so they don't remain as misleading audit leftovers.

794. assistant (2026-07-02T15:35:11.644Z)

The remaining open audit PR is khelsutra #141. Current master still has a few present-tense droplet/live claims, and #141 conflicts, so I'm re-cutting that docs fix from current `origin/master` rather than merging the stale branch.

795. assistant (2026-07-02T15:35:26.383Z)

The cherry-pick conflicts in the expected docs: `NEXT\_STEPS.md`, `DEPLOY\_ALPHA.md`, and `GPU\_VM\_SERVING.md`. I'm resolving those by preserving current master's later wording while applying the retirement/owner-gated cleanup from #141.

796. assistant (2026-07-02T15:35:31.091Z)

Most conflict choices are straightforward: keep the richer ?single-origin default / CF cleanup pending? wording from #141, plus current master's later CI and product-status wording. I'm opening the GPU runbook conflict next, then I'll patch all three conflicted files.

797. assistant (2026-07-02T15:35:37.548Z)

I've got the conflict contents. I'll patch the three conflicted files directly: the final wording will avoid ?live droplet? claims, keep Cloudflare as canonical, and preserve the newer Workers/CI wording from master.

798. assistant (2026-07-02T15:36:04.327Z)

`NEXT\_STEPS.md` is resolved. I'm applying the same conflict cleanup to the alpha deploy runbook now.

799. assistant (2026-07-02T15:36:33.909Z)

The alpha deploy runbook is resolved with the single-origin decision and current Python/GPU wording. I'm resolving the GPU VM runbook conflicts next.

800. assistant (2026-07-02T15:36:58.476Z)

The three conflict files are patched. I'm scanning for any leftover conflict markers and present-tense droplet claims, then I'll make the small follow-up doc wording fixes that the scan catches.

801. assistant (2026-07-02T15:37:09.126Z)

No conflict markers remain. The scan still found a few true stale phrases in `LOCAL\_APPLIANCE\_ARCHITECTURE.md` and one rough line wrap in `NEXT\_STEPS.md`; I'm tightening those before marking the conflict resolved.

802. assistant (2026-07-02T15:37:15.420Z)

I'm also cleaning the portability doc's example so it doesn't suggest the retired droplet key as the default key path. It's small, but it's the sort of stale breadcrumb that causes future confusion.

803. assistant (2026-07-02T15:37:49.938Z)

The stale phrasing cleanup is applied. I'm rescanning now; remaining droplet mentions should be explicitly historical or owner-actionable, not live-state claims.

804. assistant (2026-07-02T15:38:01.468Z)

The scan is now down to explicit historical/retired references and one owner-actionable SSH-key note, which is what we want. I'm checking the repo's package scripts so I can run the same local gates this docs-only PR should satisfy.

805. assistant (2026-07-02T15:38:15.657Z)

The repo's CI gates are site tests plus web typecheck/coverage/e2e/layout. Since this docs PR doesn't touch UI source, I'll still run the core local gates that prove no coverage regression: site tests, web typecheck, web coverage, plus diff whitespace.

806. assistant (2026-07-02T15:38:40.790Z)

The core local gates passed. The coverage run output was huge, so I'm reading the generated coverage summary file directly and then I'll mark the cherry-pick resolved.

807. assistant (2026-07-02T15:38:44.487Z)

Vitest didn't emit the JSON summary at that path, but the coverage command exited green, which means the configured thresholds held. I'm going to stage the resolved files explicitly to finish the cherry-pick.

808. assistant (2026-07-02T15:38:53.872Z)

Before committing, I'm also running the web Playwright/contract smoke. This PR is docs-only, but it's cheap enough here and matches the repo's CI shape more closely.

809. assistant (2026-07-02T15:39:53.623Z)

The Playwright/contract smoke is green: build passed, contract smoke passed, and all 5 CUIs passed. I'm checking status for generated artifacts, then I'll finish the cherry-pick commit with explicit paths only.

810. assistant (2026-07-02T15:39:58.897Z)

There are no stray generated files; only the intended docs are changed. I'm staging those exact doc paths to mark the cherry-pick resolved.

811. assistant (2026-07-02T15:40:05.832Z)

The index has exactly the seven docs. I'm completing the cherry-pick now, then I'll push this as the replacement PR for stale #141.

812. assistant (2026-07-02T15:40:21.604Z)

The replacement commit is ready on a fresh branch. I'm going to open a new PR rather than force-push over the stale #141 branch, then close #141 as superseded once the new PR is merged.

813. assistant (2026-07-02T15:40:24.692Z)

The replacement branch has one commit and clean whitespace. I'm pushing it and opening the PR with the local coverage/test evidence.

814. assistant (2026-07-02T15:40:30.541Z)

The branch is up. Opening the PR now; after CI is green I'll merge it and close the stale conflicting #141.

815. assistant (2026-07-02T15:40:40.288Z)

Khelsutra PR #150 is open. I'm watching its CI; if green, I'll merge it and then close #141 as superseded.

816. user (2026-07-02T15:42:52.004Z)

please close any stale PRs if they have been completely subsumed by future work

817. assistant (2026-07-02T15:43:06.588Z)

Yep, exactly. #56 was worth merging because it carried unique small Roora golden-label fixtures; #54/#55 were stale and have already been closed as superseded. I'm merging the fresh replacement for #141 now, then I'll close #141 and do a repo-wide open-PR sweep.

818. assistant (2026-07-02T15:43:17.050Z)

#150 is merged to khelsutra master as `1639804`. I'm closing the old conflicting #141 with a supersession note now.

819. assistant (2026-07-02T15:43:29.568Z)

The stale khelsutra PR is closed. I'm checking all touched repos for any remaining open PRs so we don't leave old audit branches hanging around.

820. assistant (2026-07-02T15:43:36.371Z)

Open PR sweep is clean for engine, khelsutra, rally-annotator, and sports-data-collector. I'm checking `sports-obsreport` too, then I'll refresh the umbrella tracker to include the late #56/#150 cleanup.

821. assistant (2026-07-02T15:43:45.259Z)

All five repos have zero open PRs now. I'm cutting a tiny engine docs branch from `origin/master` so the umbrella tracker records the late merges/closures rather than ending at the earlier #439 state.

822. assistant (2026-07-02T15:43:58.605Z)

The docs branch is fresh from engine `origin/master`. I'm locating the exact tracker/status lines to update: repo heads, P27/P26 notes, changelog, and data doc references.

823. assistant (2026-07-02T15:44:15.636Z)

The tracker still accurately says P0-P28 closed, but it predates the final stale-PR cleanup: sports-data #56 and khelsutra #150. I'm adding those as late cleanup records and updating the repo-head table.

824. assistant (2026-07-02T15:44:27.726Z)

I'm updating the live tracker in two places: the audit doc and the GCS data map. The changes are small: late repo heads, stale PR dispositions, sports-data #56, and khelsutra #150.

825. assistant (2026-07-02T15:44:57.833Z)

The audit doc now records the late cleanups. I'm patching `DATA\_IN\_GCS.md` so the Roora/golden-label data map matches sports-data #56 too.

826. assistant (2026-07-02T15:45:23.572Z)

The data map now includes the Roora distilled goldens and points at sports-data #56. I'm running the docs checks and stale-marker scan before committing.

827. assistant (2026-07-02T15:45:37.255Z)

Docs checks are clean. The only `this PR` marker left is intentionally the new late-cleanup changelog entry; I'll replace it with the PR number after opening the PR, just like before. Now I'm doing the branch/HEAD pre-stage check.

828. assistant (2026-07-02T15:45:41.547Z)

The engine branch is fresh from `c30ff22`, no commits above origin yet, and open PR lists are empty for the main affected repos. I'm staging only the two docs and committing the late cleanup reconciliation.

829. assistant (2026-07-02T15:45:49.157Z)

The late cleanup commit is `cce3fa1`. I'm checking that it's the only commit above `origin/master`, then I'll push and open the final docs PR.

830. assistant (2026-07-02T15:45:51.814Z)

Single-commit branch, clean tree. Pushing and opening the PR now.

831. assistant (2026-07-02T15:45:58.371Z)

The branch is pushed. I'm opening the PR with the exact verification and stale-PR closure notes.

832. assistant (2026-07-02T15:46:07.708Z)

PR #440 is open. I'm replacing the last `this PR` marker with `#440`, then I'll amend, push, and watch CI.

833. assistant (2026-07-02T15:46:19.156Z)

The PR number is now embedded in the tracker. I'm re-running docs checks and amending the branch with that final polish.

834. assistant (2026-07-02T15:46:22.278Z)

Checks are clean and the branch still has just the one docs commit. I'm staging the audit doc, amending, and force-with-lease pushing the PR branch.

835. assistant (2026-07-02T15:46:29.163Z)

PR #440 is updated. I'm watching CI; after it's green I'll merge and refresh the memory files one final time.

836. assistant (2026-07-02T15:46:37.459Z)

Checks haven't attached yet, so I'm reading the workflow run list directly like before.

837. assistant (2026-07-02T15:46:42.763Z)

The PR workflow is running. I'll watch the latest CI run through completion.

838. assistant (2026-07-02T15:49:50.023Z)

#440 CI is green, including the full offline suite with coverage floor, ruff, mypy, wheel smoke, link check, and cross-OS golden fixtures. I'm merging it now.

839. assistant (2026-07-02T15:50:09.026Z)

#440 is merged to engine master as `204dd00`. I'm refreshing the final handoff/current-state and the root pointer now so they match the actual final heads after the late cleanups.

840. assistant (2026-07-02T15:50:14.001Z)

Engine, khelsutra, sports-data, and rally-annotator are clean with no open PRs. I'm checking obsreport and then patching the shared memory files.

841. assistant (2026-07-02T15:51:06.015Z)

The memory files and root pointer now reflect the real final state: engine `204dd00`, khelsutra `1639804`, sports-data `1423e84`, and no open remediation PRs. I'm doing the final status sweep across all repos.

842. assistant (2026-07-02T15:51:11.159Z)

All five git repos are clean on their default branches, and each open-PR list is empty. I'm marking the work complete and then I'll give you the final state with the leftovers table.

843. assistant (2026-07-02T15:51:25.451Z)

Completed. The audit remediation is now fully closed out.

Merged/closed final items:

- Engine `#440` merged: final tracker cleanup, head `204dd00`.
- Khelsutra `#150` merged: retired-droplet docs refreshed, head `1639804`.
- Sports-data `#56` merged: Roora pilot golden-label fixtures, head `1423e84`.
- Closed stale/subsumed PRs: `sports-data-collector#54`, `sports-data-collector#55`, `khelsutra#141`.
- All five repos have no open PRs and clean working trees.

Docs refreshed:

- Live tracker: [AUDIT_REPORT_khelsutra-guru_2026-07-02.md](C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/docs/AUDIT_REPORT_khelsutra-guru_2026-07-02.md)
- Data map: [DATA_IN_GCS.md](C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer/docs/DATA_IN_GCS.md)
- Root pointer and memory files updated:
 - [AUDIT_REPORT_khelsutra-guru_2026-07-02.md](C:/Users/avidu/Projects/khelsutra-guru/AUDIT_REPORT_khelsutra-guru_2026-07-02.md)
 - [session-handoff.md](C:/Users/avidu/OneDrive/Documents/claude-sync/memory/khelsutra-guru/session-handoff.md)
 - [current-state.md](C:/Users/avidu/OneDrive/Documents/claude-sync/memory/khelsutra-guru/current-state.md)

Remaining leftovers are owner-gated/future backlog only:

Area	Leftover	Status
Corpus videos	Decide/upload large source videos: full 4K originals vs proxies, plus wider archive scope.	Owner-gated
GCS eval/litmus	Make eval/litmus consume `gs://` data directly.	Future engineering
Hosted deployment	Enable upload retention in live infra and finish Cloudflare/alpha cleanup.	
Owner infra/timing		
Product backlog	Non-audit feature/quality issues.	Normal backlog

Verification highlights: engine `#440` CI green including full offline suite with coverage floor, ruff, mpy, wheel smoke, link check, and cross-OS golden fixtures; khelsutra `#150` CI green plus local site/web/e2e checks; sports-data `#56` local full suite `293` passed, `2` skipped, coverage `82.02%` against floor `80`, CI green.

844. user (2026-07-02T15:51:42.507Z)

very well done!

845. assistant (2026-07-02T15:51:48.483Z)

Thank you. That was a proper marathon, but the board is finally clean and the tracker is telling the truth.