

codex-windows / 019f2174-3b5a-7980-940c-f971c36d41a7

Metadata

```
- Source: `codex-windows`
- Kind: `codex`
- Source file: `C:\Users\avidu\codex\archived_sessions\rollout-2026-07-02T11-41-33-019f2174-3b5a-7980-940c-f971c36d41a7.jsonl`
- SHA-256: `5e367e9cfd67fecc4d9961a4b7a05574301d965a95cecb9b2aa75b37d36f5acf`
- Source modified: `2026-07-02T06:11:33+00:00`
- Imported at: `2026-07-05T16:48:16+00:00`
- cli_version: `0.142.5`
- cwd: `C:\Users\avidu\Projects\badminton-highlight-indexer`
- model_provider: `openai`
- session_id: `019f2174-3b5a-7980-940c-f971c36d41a7`
- source: `vscode`
```

Transcript

1. user (2026-07-02T06:11:33.340Z)

Improve `run.log` readability for the owner's manual-QA workflow: he runs the tool on a video, then eyeballs the rally CSV against the log to verify quality. Two changes, design-then-implement, ONE PR off master.

FIRST ? branching safety (this checkout is SHARED by concurrent sessions): `git fetch origin` then `git checkout -b <branch> origin/master` (branch from the REMOTE, not the current branch). Re-verify `git branch --show-current` + `git log --oneline -1` right before committing; before the PR, `git log --oneline origin/master..HEAD` must show ONLY your commits. Stage EXPLICIT paths (never `git add -A`). Don't merge without owner.

INVESTIGATE (don't assume): find where logging is configured (`grep -rn "run.log\|basicConfig\|FileHandler\|getLogger" backend/ run\_\*.py`), how the Gemini path logs (it uses `google-genai` ? `httpx`/`grpc`/`google.auth` ? those emit the noisy per-request HTTP lines), where per-rally lines are already logged (`[rally]` / `[<label> rally]` in `backend/pipeline/segmenters/{trajectory\_hybrid,yolo\_hybrid}.py`), and where the rally CSV is written (`backend/eval/export\_wasb\_segments.py::write\_segments\_csv` and/or the reporting layer `backend/reporting/`). Read `docs/CODE\_MAP.md` for the map.

CHANGE 1 ? suppress HTTP chatter: raise the noisy third-party loggers (`httpx`, `httpcore`, `urllib3`, `google\_genai`/`google.genai`, `google.auth`, `grpc`) to WARNING by default, behind a config knob (e.g. `logging.verbose\_http: false`) so debugging can re-enable them. Do it in the central logging setup, not scattered. Don't suppress our own `backend.\*` logs.

CHANGE 2 ? aggregate rally-times summary: at the END of a run, log a compact, scannable block ? total rallies, total highlight duration, and a per-rally table `# start(HH:MM:SS) end dur(s) shot\_count` ? sourced from the produced play_segments (the same data that goes to the rally CSV), so the log and CSV line up for manual verification. Make it a small pure helper (format a list of segments ? text) so it's unit-testable offline, then call it from the run path. Keep timestamps in HH:MM:SS to match the existing `\_fmt\_ts`.

Deliverable: a short plan in the PR description + the implementation + tests (`python run\_all\_tests.py` green, ruff + mypy clean). Keep it minimal and config-gated; default behavior = quieter log + the summary block. Don't touch unrelated areas.

2. assistant (2026-07-02T06:11:33.340Z)

I'll start by setting up the branch safely and investigating the codebase in parallel.

3. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_call: Bash]
description: Fetch remote and show origin/master tip
```

```
command: git fetch origin && git log --oneline -3 origin/master
[/external_agent_tool_call]
```

4. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_call: Grep]
input: {"pattern":"run\\.log|basicConfig|FileHandler|getLogger|setLevel","path":"backend","output_mode":"content","-n":true}
[/external_agent_tool_call]
```

5. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_call: Grep]
input: {"pattern":"run\\.log|basicConfig|FileHandler|getLogger|setLevel|setup_logging|configure_logging","glob":"run_*.py","output_mode":"content","-n":true}
[/external_agent_tool_call]
```

6. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_result]
35a50f3 docs(readme): index QUALITY_ITERATIONS + mark EXPERIMENT_HARNESS P1 SHIPPED (#118)
cbf2236 feat(segmenters): FusionSegmenter ? shuttle + net-crossing + person-cue fusion (P2, default-OFF) (#117)
df7d5df refactor(detectors): extract the person cue into PersonDetector (fusion P1) (#115)
[/external_agent_tool_result]
```

7. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_result]
run_regression.py:53: logging.basicConfig(level=logging.INFO, format="%(levelname)s %(message)s")
run_eval.py:75: logging.basicConfig(level=logging.INFO, format="%(levelname)s %(message)s")
run_phase3_eval.py:41: logger = logging.getLogger(__name__)
run_phase3_eval.py:76: logging.basicConfig(level=logging.INFO, format="%(levelname)s %(message)s")
run_process_and_stitch.py:8: # Configure root logger ? all backend modules use
logging.getLogger(__name__)
run_process_and_stitch.py:9: logging.basicConfig(
scratch\run_gx010125.py:18: logging.basicConfig(
[/external_agent_tool_result]
```

8. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_result]
backend\main.py:15: logging.basicConfig(
backend\main.py:20: logger = logging.getLogger(__name__)
backend\config\loader.py:20: logger = logging.getLogger(__name__)
backend\providers\gemini.py:11: logger = logging.getLogger(__name__)
backend\infrastructure\database.py:7: logger = logging.getLogger(__name__)
backend\utils\video.py:5: logger = logging.getLogger(__name__)
backend\annotations\automated_provider.py:30: logger = logging.getLogger(__name__)
backend\eval\gt_loader.py:26: logger = logging.getLogger(__name__)
backend\pipeline\stitcher\compiler.py:8: logger = logging.getLogger(__name__)
backend\eval\experiment.py:55: logger = logging.getLogger(__name__)
backend\pipeline\segmenters\gemini.py:13: logger = logging.getLogger(__name__)
backend\pipeline\detectors\person_detector.py:33: logger = logging.getLogger(__name__)
backend\pipeline\segmenters\motion.py:9: logger = logging.getLogger(__name__)
backend\pipeline\detectors\base.py:75: logger = logging.getLogger(__name__)
backend\eval\harness.py:25: logger = logging.getLogger(__name__)
backend\eval\regression.py:29: logger = logging.getLogger(__name__)
backend\pipeline\segmenters\trajectory_hybrid.py:34: logger = logging.getLogger(__name__)
backend\pipeline\detectors\tracknet_runner.py:27: logger = logging.getLogger(__name__)
backend\pipeline\segmenters\yolo_hybrid.py:30: logger = logging.getLogger(__name__)
backend\eval\annotations.py:19: logger = logging.getLogger(__name__)
```

```
backend\pipeline\detectors\wasb_runner.py:24:logger = logging.getLogger(__name__)
backend\pipeline\detectors\wasb_infer.py:48:logger = logging.getLogger(__name__)
[/external_agent_tool_result]
```

9. assistant (2026-07-02T06:11:33.340Z)

Let me look at where `run.log` is configured and read the central logging setup files.

10. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_call: Grep]
input: {"pattern":"run\\.log|\\.log[\\"]|FileHandler|verbose_http|setup_logging|configure_loggin
g","output_mode":"content","-n":true,"head_limit":60}
[/external_agent_tool_call]
```

11. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_result]
tools\cleanup_caches.py:58:     if low.endswith(("log", ".sh", ".txt", ".json", ".md")):
[/external_agent_tool_result]
```

12. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\run_process_and_stitch.py
[/external_agent_tool_call]
```

13. assistant (2026-07-02T06:11:33.340Z)

```
[external_agent_tool_result]
1     import os
2     import json
3     import logging
4     from dotenv import load_dotenv
5
6     load_dotenv()
7
8     # Configure root logger ? all backend modules use logging.getLogger(__name__)
9     logging.basicConfig(
10         level=logging.INFO,
11         format="%(asctime)s [%(name)s] %(levelname)s: %(message)s",
12         datefmt="%H:%M:%S",
13     )
14
15     from backend.infrastructure.database import Database
16     from backend.config import load_config
17     from backend.utils.hashing import compute_video_id
18     from backend.pipeline.segmenters import segmenter_registry
19     from backend.pipeline.stitcher.compiler import VideoCompiler
20
21     def main():
22         video_path = r"C:\Users\avidu\Videos\Testing\KushagraYashAviFirstVideo.mp4"
23         output_dir = r"C:\Users\avidu\Videos\Testing"
24         output_filename = "Final_2_OutputOfKushagraYashAviFirstVideo.mp4"
25
26         print("Loading config...")
27         cfg = load_config() # typed AppConfig ? single source of defaults
28
29         print("Verifying files exist...")
30         if not os.path.exists(video_path):
31             print(f"Error: Video file not found at {video_path}")
32             return
33
34         os.makedirs(cfg.output.output_dir, exist_ok=True)
```

```

35     db_path = os.path.join(cfg.output.output_dir, cfg.output.db_name)
36     db = Database(db_path)
37
38     # Calculate MD5 of the video to uniquely identify it
39     import time
40     print("Calculating video MD5 (this might take a few seconds on large files)...",
flush=True)
41     t0 = time.time()
42     video_id = compute_video_id(video_path)
43     print(f"Calculated MD5: {video_id} (took {time.time()-t0:.1f}s)", flush=True)
44
45     segments = []
46     cached_video = db.get_video(video_id)
47     cache_hit = False
48
49     if cached_video and cached_video.get("status") == "processed":
50         segments = db.query_play_segments(video_id, {})
51         if len(segments) > 0:
52             print(f"\n[CACHE] Found cached indexed video in database (MD5: {video_id})
with {len(segments)} parsed play segments.", flush=True)
53             choice = input("Would you like to (1) Run stitching on the cached segments,
or (2) Reprocess the video with Gemini from scratch? [1/2]: ").strip()
54             if choice == '1':
55                 print("[CACHE] Using cached play segments. Skipping segmenter API
call.", flush=True)
56                 cache_hit = True
57             else:
58                 print("[CACHE] Ignoring cache...", flush=True)
59             else:
60                 print("[CACHE] Video exists in DB but contains 0 play segments. Treating as
cache miss.", flush=True)
61
62         if not cache_hit:
63             print("[CACHE] Initializing clean indexing...", flush=True)
64             db.clear_video_data(video_id)
65
66         print("Bypassing OpenCV validations for performance on massive files...",
flush=True)
67         db.add_video(video_id, video_path, "processed", [])
68
69         print("Running segmenter & indexing...", flush=True)
70         seg_name = cfg.indexing.default_segmenter
71         segmenter_cls = segmenter_registry.get(seg_name)
72         if not segmenter_cls:
73             print(f"[FAIL] Segmenter '{seg_name}' not found.")
74             return
75         print(f"Using segmenter: {seg_name}", flush=True)
76         segmenter = segmenter_cls(db, cfg)
77         segments = segmenter.process_video(video_id, video_path)
78         print(f"Successfully indexed {len(segments)} play segments.", flush=True)
79
80     if not segments:
81         print("No play segments detected. Cannot compile highlights.")
82         return
83
84
85     end_padding = cfg.stitching.end_padding
86     begin_padding = cfg.stitching.begin_padding
87     min_shot_count = cfg.stitching.min_shot_count
88
89     if cache_hit:
90         print("\n--- Stitching Customizations ---")
91         try:
92             bpad_input = input(f"Enter begin_padding in seconds (default
{begin_padding}): ").strip()

```

```
93             if bpad_input:
94                 begin_padding ...
[/external_agent_tool_result]
```

14. assistant (2026-07-02T06:11:33.340Z)

`run.log` isn't a literal in code ? it's a stdout/stderr redirect capturing the `logging` output configured by `basicConfig`. Let me read the central pieces: `main.py` logging, the config structure, the segmenter rally-logging, and the CSV writer.

15. assistant (2026-07-02T06:11:33.341Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\main.py
[/external_agent_tool_call]
```

16. assistant (2026-07-02T06:11:33.341Z)

```
[external_agent_tool_result]
1     import os
2     import sys
3     import argparse
4     import logging
5     import threading
6     import traceback
7     import uuid
8     from concurrent.futures import ThreadPoolExecutor
9     import uvicorn
10    from dotenv import load_dotenv
11
12    load_dotenv()
13
14    # Centralized logging setup for the indexer (hot paths use logger, CLI output uses print
for direct user-facing messages).
15    logging.basicConfig(
16        level=logging.INFO,
17        format='%(asctime)s [%(levelname)s] %(name)s: %(message)s',
18        datefmt='%H:%M:%S'
19    )
20    logger = logging.getLogger(__name__)
21    from fastapi import FastAPI, HTTPException, Request
22    from fastapi.middleware.cors import CORSMiddleware
23    from pydantic import BaseModel
24    from typing import List, Dict, Any, Optional
25
26    from fastapi.staticfiles import StaticFiles
27    from fastapi.responses import FileResponse, HTMLResponse, JSONResponse
28    from backend.pipeline.validators.base import registry
29    from backend.infrastructure.database import Database
30    from backend.pipeline.segmenters.base import segmenter_registry
[/external_agent_tool_result]
```

17. assistant (2026-07-02T06:11:33.341Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\eval\export_wasb_segments.py
[/external_agent_tool_call]
```

18. assistant (2026-07-02T06:11:33.341Z)

```
[external_agent_tool_call: Glob]
input: {"pattern":"backend/config/**/*.py"}
[/external_agent_tool_call]
```

19. assistant (2026-07-02T06:11:33.341Z)

```
[external_agent_tool_result]
1      """Export tuned WASB rally segments for manual eyeball verification.
2
3      After calibration picks a tuned threshold config, this turns a cached WASB
4      trajectory into the actual rally [start,end] segments that config produces, so the
5      user can answer "are the rallies actually better?" ? by scrubbing the video to each
6      segment, by diffing tuned-vs-golden quantitatively, or by cutting clips.
7
8      Reuses (no new windowing/metrics/cutting logic):
9      - calibrate_wasb.make_predict_fn -> trajectory -> segments for a given config
10     - eval.metrics.match_intervals/score -> tuned-vs-golden IoU comparison
11     - tools.rally_slicer.slice_rallies -> optional clip cutting (the exported CSV is
12       written in the golden/slicer schema, so it feeds straight back in)
13
14     Run (uses the tuned config from the first calibration by default):
15     python -m backend.eval.export_wasb_segments --trajectory
~/clips/sample_long_traj.csv \
16         --labels "Annotation Setup/golden_labels_badminton.csv" --fps 59.94 --frame-
width 1920
17         # add --video <mp4> --slice to also cut one padded clip per tuned rally
18     """
19     from __future__ import annotations
20
21     import csv
22     import os
23     import os.path as osp
24     import argparse
25     from typing import List, Tuple, Dict, Optional
26
27     from backend.eval.calibrate_wasb import make_predict_fn, load_golden_gts, BASELINE
28     from backend.eval.metrics import match_intervals, score
29     from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
30     from backend.pipeline.detectors import rally_gate
31
32     Interval = Tuple[float, float]
33     Config = Tuple[float, float, float]
34
35     # The tuned config from the first per-video calibration (velocity, merge_gap, min_dur);
36     # F1 0.588 (baseline) -> 0.746. Override on the CLI to export a different operating
point.
37     TUNED: Config = (0.05, 1.0, 1.0)
38
39     SEG_FIELDS = ["rally_number", "start_time", "end_time", "duration", "start_mmss",
"end_mmss"]
40     COMP_FIELDS = ["type", "ref", "start_time", "end_time", "mmss", "status",
41                   "iou", "pred_start", "pred_end", "start_err_s", "end_err_s"]
42
43
44     def seconds_to_mmss(t: float) -> str:
45         """Human-scrubbable timestamp, e.g. 106.031 -> '1:46.031'."""
46         t = max(0.0, float(t))
47         m = int(t // 60)
48         return f"{m}:{t - m * 60:06.3f}"
49
50
51     def segments_to_rows(segments: List[Interval]) -> List[Dict[str, str]]:
52         """Number segments 1..N in time order; emit the golden/slicer CSV schema."""
53         rows = []
54         for i, (s, e) in enumerate(segments, start=1):
55             rows.append({
56                 "rally_number": str(i),
57                 "start_time": f"{s:.3f}",
58                 "end_time": f"{e:.3f}",
```

```

59         "duration": f"{e - s:.3f}",
60         "start_mmss": seconds_to_mmss(s),
61         "end_mmss": seconds_to_mmss(e),
62     })
63     return rows
64
65
66     def _ensure_parent(path: str) -> None:
67         os.makedirs(osp.dirname(osp.abspath(path)) or ".", exist_ok=True)
68
69
70     def write_segments_csv(path: str, segments: List[Interval]) -> str:
71         """Write segments in the golden schema (rally_number/start_time/end_time + mm:ss).
72
73         The schema is what tools.rally_slicer.load_rally_labels consumes, so the file
74         doubles as input to clip cutting.
75         """
76         _ensure_parent(path)
77         with open(path, "w", newline="", encoding="utf-8") as f:
78             w = csv.DictWriter(f, fieldnames=SEG_FIELDS)
79             w.writeheader()
80             w.writerows(segments_to_rows(segments))
81         return path
82
83
84     def build_comparison(preds: List[Interval], gts: List[Interval],
85                         iou_threshold: float = 0.5) -> dict:
86         """Golden-centric comparison: per golden rally, did a tuned segment match (IoU),
87         plus the boundary errors and any extra (false-positive) tuned segments."""
88         mr = match_intervals(preds, gts, iou_threshold)
89         sc = score(preds, gts, iou_threshold, match_result=mr)
90         ...
[/external_agent_tool_result]

```

20. assistant (2026-07-02T06:11:33.341Z)

```

[/external_agent_tool_result]
backend\config\loader.py
backend\config\__init__.py
backend\config\models.py
[/external_agent_tool_result]

```

21. assistant (2026-07-02T06:11:33.341Z)

The WASB export CSV is the eval path. The owner's manual-QA runs `run\_process\_and\_stitch.py` over play_segments. Let me find `\_fmt\_ts`, the reporting/CSV layer, the rally-logging lines, and the config model.

22. assistant (2026-07-02T06:11:33.341Z)

```

[/external_agent_tool_call: Grep]
input: {"pattern": "_fmt_ts|def _fmt_ts|\\[rally\\]|rally\\]|shot_count", "path": "backend", "output_mode": "content", "-n": true, "head_limit": 50}
[/external_agent_tool_call]

```

23. assistant (2026-07-02T06:11:33.341Z)

```

[/external_agent_tool_result]
backend\config\models.py:193:     min_shot_count: int = 1
backend\main.py:502:     # The configured shot-count floor (stitching.min_shot_count) applies on
the API
backend\main.py:503:     # path too ? parity with run_process_and_stitch.py. Segments WITHOUT
shot_count
backend\main.py:510:     sc = meta.get("shot_count") if isinstance(meta, dict) else None

```

```

backend\main.py:512:         if sc is not None and int(sc) < stitching.min_shot_count:
backend\main.py:518:         logger.info(f"[compile]
stitching.min_shot_count={stitching.min_shot_count}: dropped "
backend\main.py:523:
f"stitching.min_shot_count={stitching.min_shot_count}.")
backend\tools\export.py:76:         """Return a CSV string
(start,end,duration,shot_count,detector,...) .
backend\tools\export.py:84:
lines.append("start_time,end_time,duration,shot_count,shot_count_confidence,detector,id")
backend\tools\export.py:87:         shot = md.get("shot_count", "")
backend\tools\export.py:88:         conf = md.get("shot_count_confidence", "")
backend\sports\badminton.py:35:Confidence Scoring: Because the shuttlecock moves at extreme
velocities that may fall between sampled frames, provide a shot_count_confidence metric (Low,
Medium, High). Base your estimate on player biomechanics, positioning, and racket swings if the
shuttle itself is blurred.
backend\sports\badminton.py:46:"shot_count_confidence": "High"
backend\sports\badminton.py:52:"shot_count_confidence": "Medium"
backend\tools\rally_slicer.py:72:         labels[rally] = (start, end)
backend\tools\rally_slicer.py:93:         start, end = labels[rally]
backend\eval\harness.py:85:def _fmt_ts(seconds: float) -> str:
backend\eval\harness.py:134:         lines.append(f" -
{_fmt_ts(iv[0])}?{_fmt_ts(iv[1])}")
backend\eval\harness.py:137:         lines.append(f" -
{_fmt_ts(iv[0])}?{_fmt_ts(iv[1])}")
backend\pipeline\segmenters\gemini.py:428:         # Extract shot_count from Gemini or
compute estimate if not present
backend\pipeline\segmenters\gemini.py:429:         shot_count =
segment.get("shuttle_exchange_count")
backend\pipeline\segmenters\gemini.py:430:         if shot_count is None:
backend\pipeline\segmenters\gemini.py:431:             shot_count = max(3, int(dur / 0.8)) #
1 shot every 0.8 seconds as fallback estimate
backend\pipeline\segmenters\gemini.py:434:             shot_count = int(shot_count)
backend\pipeline\segmenters\gemini.py:436:             shot_count = max(3, int(dur /
0.8))
backend\pipeline\segmenters\gemini.py:440:             "shot_count": shot_count,
backend\pipeline\segmenters\gemini.py:441:             "shot_count_confidence":
segment.get("shot_count_confidence", "Unknown"),
backend\pipeline\segmenters\motion.py:168:             "shot_count": random.randint(4, 25),
backend\pipeline\segmenters\trajectory_hybrid.py:37:def _fmt_ts(seconds: float) -> str:
backend\pipeline\segmenters\trajectory_hybrid.py:67:     def _shot_count_from(segment: Dict[str,
Any], duration: float) -> int:
backend\pipeline\segmenters\trajectory_hybrid.py:73:         shot_count =
segment.get("shuttle_exchange_count")
backend\pipeline\segmenters\trajectory_hybrid.py:74:         if shot_count is None:
backend\pipeline\segmenters\trajectory_hybrid.py:77:             return int(shot_count)
backend\pipeline\segmenters\trajectory_hybrid.py:196:             logger.info(f"[{label}
rally] {_fmt_ts(cw['start'])}-{_fmt_ts(cw['end'])} "
backend\pipeline\segmenters\trajectory_hybrid.py:276:             "shot_count":
self._shot_count_from(segment, dur),
backend\pipeline\segmenters\trajectory_hybrid.py:277:
"shot_count_confidence": segment.get("shot_count_confidence", "Unknown"),
backend\pipeline\segmenters\yolo_hybrid.py:33:def _fmt_ts(seconds: float) -> str:
backend\pipeline\segmenters\yolo_hybrid.py:241:         f"[rally]
{_fmt_ts(w['start'])}-{_fmt_ts(w['end'])} "
backend\pipeline\segmenters\yolo_hybrid.py:333: ...
[/external_agent_tool_result]

```

24. assistant (2026-07-02T06:11:33.341Z)

```

[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\config\models.py
[/external_agent_tool_call]

```

25. assistant (2026-07-02T06:11:33.341Z)


```
[external_agent_tool_call: Glob]
input: {"pattern":"backend/reporting/**"}
[/external_agent_tool_call]
```

26. assistant (2026-07-02T06:11:33.341Z)

```
[external_agent_tool_result]
1      """Pydantic models for the application configuration.
2
3      This module defines the canonical, validated, typed configuration for the
4      Badminton Highlight Indexer. Built-in defaults live here (single source of truth).
5
6      The models support graceful migration from old config.json files via extra="allow".
7      """
8
9      from __future__ import annotations
10
11     from typing import Any, List, Literal, Optional
12
13     from pydantic import BaseModel, Field, field_validator, model_validator
14
15     Sport = Literal["badminton", "tennis", "pickleball", "table_tennis", "cricket"]
16
17
18     class ValidationRelevanceConfig(BaseModel):
19         court_green_lower: list[int] = Field(default=[35, 40, 40])
20         court_green_upper: list[int] = Field(default=[85, 255, 255])
21         court_pixels_min_ratio: float = Field(default=0.05, ge=0.0, le=1.0)
22
23
24     class ValidationConfig(BaseModel):
25         min_resolution_width: int = 1280
26         min_resolution_height: int = 720
27         min_fps: float = 29.9
28         min_blur_threshold: float = 20.0
29         max_scene_cuts_per_minute: float = 0.5
30         relevance: ValidationRelevanceConfig =
31         Field(default_factory=ValidationRelevanceConfig)
32
33     class TrackNetConfig(BaseModel):
34         wsl_distro: str = "Ubuntu"
35         repo_dir: str = "~/models/TrackNetV4"
36         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
37         conda_env: str = "TrackNetV4"
38         weights_path: str = ""
39         queue_length: int = 5
40         stage_in_wsl: bool = True
41         wsl_stage_dir: str = "~/clips"
42         timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
43         velocity_threshold: float = 0.02
44         # Cache hygiene: after a SUCCESSFUL run the staged video copy (and, for WASB, the
45         # decoded frame cache) are deleted automatically ? they are pure intermediates,
46         # regenerable from the source, and the dominant disk consumer (~GBs/video). Set
47         # true only for debugging. On FAILURE they are always kept so a re-run can resume.
48         keep_frames: bool = False
49         # Warn before a run if WSL free space (at wsl_stage_dir) is below this many GB.
50         # 0 = disabled.
51         min_free_gb: float = 0.0
52
53
54     class WasbIndexConfig(BaseModel):
55         """Typed config for the live WASB shuttle substrate (indexing.wasb).
56
57         Mirrors backend/pipeline/detectors/wasb_runner.py::WasbConfig.from_indexing_cfg so
```

```

58         these knobs actually take effect ? previously this whole block was silently dropped
59         because nested config sections defaulted to extra='ignore'.
60         """
61         wsl_distro: str = "Ubuntu"
62         repo_dir: str = "~/models/WASB-SBDT"
63         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
64         conda_env: str = "wasb"
65         weights_path: str = "~/models/WASB-
SBDT/pretrained_weights/wasb_badminton_best.pth.tar"
66         sport: str = "badminton"
67         wsl_stage_dir: str = "~/clips"
68         timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
69         velocity_threshold: float = 0.02
70         # Cache hygiene: after a SUCCESSFUL run the staged video copy + decoded frame cache
71         # (the ~GBs/video disk hog) are deleted automatically ? pure intermediates,
regenerable
72         # from the source. Set true only for debugging. On FAILURE they are kept so a re-run
resumes.
73         keep_frames: bool = False
74         # Warn before a run if WSL free space (at wsl_stage_dir) is below this many GB. 0 =
disabled.
75         min_free_gb: float = 0.0
76
77
78     class FusionConfig(BaseModel):
79         """Config for the multi-signal `fusion_hybrid` segmenter (MULTI_SIGNAL_FUSION_PLAN
P2).
80
81         Every default reproduces control behaviour: the fusion segmenter persists the
82         shuttle net-crossing count (Gate A signal) and ? only when ``cue_flags['person']``
is
83         true ? the person-cue features, but it does NOT gate the served handoff unless a
84         threshold is raised (``min_crossings``/``min_players`` default 0 = OFF). The court
mask
85         is optional...
[/external_agent_tool_result]

```

27. assistant (2026-07-02T06:11:33.341Z)

```

[/external_agent_tool_result]
backend\reporting\harvest.py
backend\reporting\run_signals.py
backend\reporting\spec.yaml
backend\reporting\__init__.py
backend\reporting\__pycache__\harvest.cpython-313.pyc
backend\reporting\__pycache__\run_signals.cpython-313.pyc
backend\reporting\__pycache__\__init__.cpython-313.pyc
[/external_agent_tool_result]

```

28. assistant (2026-07-02T06:11:33.342Z)

Let me read the reporting layer, the CSV export tool, and the `\_fmt\_ts` reference to decide the helper's home and signature.

29. assistant (2026-07-02T06:11:33.342Z)

```

[/external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\tools\export.py
[/external_agent_tool_call]

```

30. assistant (2026-07-02T06:11:33.342Z)

```

[/external_agent_tool_result]
1         """Minimal Value Layer exports (first slice of review item 3).

```

```

2
3 Provides timeline exports (EDL, JSON, CSV) and basic stats over the
4 play_segments returned by the indexer (or query_play_segments). These are
5 the "time-saved editing is table stakes" deliverables and also give a
6 machine-readable foundation for later human correction / golden feeding
7 and clip galleries.
8
9 Kept deliberately small and pure (no DB, no side effects) so they are
10 easy to call from CLI, notebooks, or a future API endpoint.
11 ""
12
13 from __future__ import annotations
14
15 from typing import Any, Dict, Iterable, List, Optional
16
17
18 def _to_list(segments: Iterable[Dict[str, Any]]) -> List[Dict[str, Any]]:
19     return list(segments) if not isinstance(segments, list) else segments
20
21
22 def basic_stats(segments: Iterable[Dict[str, Any]]) -> Dict[str, Any]:
23     """Return simple aggregate stats for a list of play segments.
24
25     Keys: count, total_duration_s, avg_duration_s, min_duration_s, max_duration_s.
26     All durations are in seconds (float).
27     """
28     segs = _to_list(segments)
29     if not segs:
30         return {
31             "count": 0,
32             "total_duration_s": 0.0,
33             "avg_duration_s": 0.0,
34             "min_duration_s": 0.0,
35             "max_duration_s": 0.0,
36         }
37     durs = [float(s.get("duration", (s.get("end_time", 0) - s.get("start_time", 0))))
38 for s in segs]
39     total = sum(durs)
40     return {
41         "count": len(segs),
42         "total_duration_s": round(total, 2),
43         "avg_duration_s": round(total / len(durs), 2),
44         "min_duration_s": round(min(durs), 2),
45         "max_duration_s": round(max(durs), 2),
46     }
47
48 def segments_to_json(
49     segments: Iterable[Dict[str, Any]],
50     video_id: str,
51     video_path: Optional[str] = None,
52     meta: Optional[Dict[str, Any]] = None,
53 ) -> Dict[str, Any]:
54     """Return a JSON-serializable timeline dict.
55
56     Shape (stable):
57     {
58         "video_id": "...",
59         "video_path": "...",
60         "meta": {...},
61         "segments": [ {"id":, "start_time":, "end_time":, "duration":, "metadata": {...}},
... ],
62         "stats": basic_stats(...)
63     }
64     """

```



```

9      store is consumed (popped) once per report, and signals for a video are simply
10     dropped if no report is emitted. No obsreport dependency ? importable everywhere.
11     """
12     from __future__ import annotations
13
14     import threading
15     from typing import Any, Dict
16
17     _lock = threading.Lock()
18     _signals: Dict[str, Dict[str, Any]] = {}
19
20
21     def incr(video_id: str, key: str, by: int = 1) -> None:
22         """Add `by` to the counter `key` for this video's in-flight run."""
23         if not video_id or by == 0:
24             return
25         with _lock:
26             sig = _signals.setdefault(video_id, {})
27             sig[key] = int(sig.get(key, 0)) + by
28
29
30     def pop(video_id: str) -> Dict[str, Any]:
31         """Consume (remove and return) all signals recorded for this video's run."""
32         with _lock:
33             return _signals.pop(video_id, {})
34
[/external_agent_tool_result]

```

33. assistant (2026-07-02T06:11:33.342Z)

```

[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-
indexer\backend\pipeline\segmenters\trajectory_hybrid.py
[/external_agent_tool_call]

```

34. assistant (2026-07-02T06:11:33.342Z)

```

[external_agent_tool_result]
30     from backend.sports import sport_registry
31     from backend.providers import provider_registry
32     from backend.pipeline.detectors.base import DetectorRunner
33
34     logger = logging.getLogger(__name__)
35
36
37     def _fmt_ts(seconds: float) -> str:
38         seconds = max(0, int(seconds))
39         h, rem = divmod(seconds, 3600)
40         m, s = divmod(rem, 60)
41         return f"{h:02d}:{m:02d}:{s:02d}"
42
43
44     class TrajectoryHybridSegmenter(BasePlaySegmenter):
45         """Trajectory-detector hybrid: precise candidate rallies + AI semantic
46         boundaries."""
47
48         #: config section under `indexing` (velocity_threshold lives there), the
49         #: output subdir for trajectories, and the metadata detector-tag prefix.
50         family: str = ""
51
[/external_agent_tool_result]

```

35. assistant (2026-07-02T06:11:33.342Z)

```

[external_agent_tool_result]

```

```
backend\utils\__init__.py
backend\utils\geometry.py
backend\utils\hashing.py
backend\utils\paths.py
backend\utils\video.py
[/external_agent_tool_result]
```

36. assistant (2026-07-02T06:11:33.342Z)

Now let me look at the test structure and confirm the `utils` package + how tests are discovered.

37. assistant (2026-07-02T06:11:33.342Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\run_all_tests.py
[/external_agent_tool_call]
```

38. assistant (2026-07-02T06:11:33.342Z)

```
[external_agent_tool_result]
1      #!/usr/bin/env python
2      """Run the FULL test suite for the badminton-highlight-indexer.
3
4      A single, discoverable entry point so you (or a reviewer / CI) can verify the
5      whole suite before opening a PR:
6
7          python run_all_tests.py          # full suite, quiet
8          python run_all_tests.py -v      # verbose (passthrough to pytest)
9          python run_all_tests.py -k config # filter by keyword (passthrough)
10         python run_all_tests.py --cov=backend --cov-report=term-missing # coverage
11
12     Anything after the script name is passed straight through to pytest.
13
14     Notes
15     -----
16     - The suite is fully mocked (no GPU, no network, no WSL) and runs offline.
17     - The observability **reporting** tests skip cleanly if the optional ``obsreport``
18       package isn't installed (they use ``pytest.importorskip``). To exercise them too:
19       pip install -e ../sports-obsreport --no-build-isolation
20     - Exit code is pytest's own, so this is CI / pre-PR friendly (0 = all green).
21     """
22     from __future__ import annotations
23
24     import importlib.util
25     import subprocess
26     import sys
27
28
29     def _installed(module: str) -> bool:
30         try:
31             return importlib.util.find_spec(module) is not None
32         except (ImportError, ValueError):
33             return False
34
35
36     def main() -> int:
37         args = list(sys.argv[1:])
38         # Default to quiet unless the caller chose a verbosity.
39         if not any(a in ("-q", "-v", "-vv", "-vvv", "--quiet", "--verbose") for a in args):
40             args = ["-q", *args]
41         cmd = [sys.executable, "-m", "pytest", *args]
42
43         obs = "available" if _installed("obsreport") else "NOT installed ? reporting tests
```

```

will SKIP"
44     bar = "=" * 72
45     print(bar)
46     print(" Full test suite: " + " ".join(cmd))
47     print(f" obsreport (optional reporting dep): {obs}")
48     print(bar, flush=True)
49
50     rc = subprocess.run(cmd).returncode
51
52     print(bar)
53     print(" RESULT: " + ("PASS - all tests green" if rc == 0 else f"FAIL - pytest exit
{rc}"))
54     print(bar)
55     return rc
56
57
58     if __name__ == "__main__":
59         raise SystemExit(main())
60
[/external_agent_tool_result]

```

39. assistant (2026-07-02T06:11:33.342Z)

```

[external_agent_tool_call: Glob]
input: {"pattern":"backend/utils/*.py"}
[/external_agent_tool_call]

```

40. assistant (2026-07-02T06:11:33.342Z)

```

[external_agent_tool_call: Glob]
input: {"pattern":"tests/**/*.export*.py"}
[/external_agent_tool_call]

```

41. assistant (2026-07-02T06:11:33.342Z)

```

[external_agent_tool_result]
tests\test_export_wasb_segments.py
[/external_agent_tool_result]

```

42. assistant (2026-07-02T06:11:33.342Z)

Let me check existing tests for the export tool and config to match test style.

43. assistant (2026-07-02T06:11:33.342Z)

```

[external_agent_tool_call: Glob]
input: {"pattern":"tests/**/*.py"}
[/external_agent_tool_call]

```

44. assistant (2026-07-02T06:11:33.343Z)

```

[external_agent_tool_result]
tests\__init__.py
tests\test_eval_batch.py
tests\test_eval_training.py
tests\test_annotation_providers.py
tests\test_calibration.py
tests\test_calibrate_wasb.py
tests\test_export_wasb_segments.py
tests\test_gemini_refine.py
tests\test_calibrate_local.py
tests\test_distill_local.py
tests\test_automated_provider.py
tests\test_eval_harness.py

```

```
tests\test_hit_detector.py
tests\test_validators_base.py
tests\test_hashing.py
tests\test_rally_slicer.py
tests\test_reporting_spec_faq.py
tests\test_tracknet_runner.py
tests\test_eval_partial_labels.py
tests\test_detector_keying.py
tests\test_reingest_safety.py
tests\test_eval_metrics.py
tests\test_compiler.py
tests\test_eval_annotations.py
tests\test_gt_loader.py
tests\test_import_smoke.py
tests\test_config_wasb.py
tests\test_reporting_harvest.py
tests\test_run_signals.py
tests\test_async_jobs.py
tests\test_base.py
tests\test_database.py
tests\test_eval_classifier.py
tests\test_eval_regression.py
tests\test_gemini.py
tests\test_geometry.py
tests\test_path_safety.py
tests\test_provider_gemini.py
tests\test_providers.py
tests\test_rally_seq_proto.py
tests\test_reliability_hardening.py
tests\test_reporting.py
tests\test_sport_profiles.py
tests\test_training_gen0.py
tests\test_video.py
tests\test_wasb_cache.py
tests\test_tracknet_hybrid.py
tests\test_wsl_tilde_quoting.py
tests\test_upload_download.py
tests\test_api_endpoints.py
tests\test_watermark.py
tests\test_cleanup_caches.py
tests\test_wasb_runner.py
tests\test_experiment.py
tests\test_yolo_hybrid.py
tests\test_config.py
tests\test_fusion_features.py
tests\test_fusion_hybrid.py
tests\test_person_detector.py
tests\test_rally_gate.py
tests\test_trajectory_hybrid_equivalence.py
[/external_agent_tool_result]
```

45. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_call: Grep]
input: {"pattern":"segments_to_csv|basic_stats|from backend.tools.export|import
export","path":"tests","output_mode":"files_with_matches"}
[/external_agent_tool_call]
```

46. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_result]
Found 2 files
tests\test_rally_slicer.py
tests\test_export_wasb_segments.py
[/external_agent_tool_result]
```


47. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\tests\test_config.py
[/external_agent_tool_call]
```

48. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_result]
1      """Basic tests for the typed configuration package.
2
3      These ensure:
4      - Defaults are sensible and come from the models.
5      - File overrides work.
6      - Old config files with extra keys are tolerated.
7      - save_default_config produces loadable output.
8      """
9
10     import json
11     import tempfile
12     import os
13
14
15     from backend.config import load_config, save_default_config, AppConfig
16
17
18     def test_load_default_config():
19         # Force load without the project's config.json so we test the model built-in
defaults
20         cfg = load_config("/nonexistent/path/that/does/not/exist/12345/config.json")
21         assert isinstance(cfg, AppConfig)
22         assert cfg.sport == "badminton"
23         assert cfg.indexing.default_segmenter == "yolo_hybrid"
24         assert cfg.validation.min_blur_threshold == 20.0 # updated default from review
25         assert cfg.indexing.ai_max_workers == 5
26
27
28     def test_fusion_config_defaults_reproduce_control():
29         """The fusion block defaults must be no-ops: wasb substrate, no cues, no gating, no
mask."""
30         cfg = load_config("/nonexistent/path/12345/config.json")
31         f = cfg.indexing.fusion
32         assert f.substrate == "wasb"
33         assert f.cue_flags == {} # person cue OFF
34         assert f.min_crossings == 0 # Gate A OFF
35         assert f.min_players == 0 # Gate B OFF
36         assert f.court_polygon is None # no mask
37         assert f.velocity_threshold == 0.02
38
39
40     def test_fusion_config_file_override():
41         with tempfile.TemporaryDirectory() as td:
42             p = os.path.join(td, "fusion.json")
43             with open(p, "w") as fh:
44                 json.dump({"indexing": {"fusion": {
45                     "substrate": "tracknet", "min_crossings": 2,
46                     "cue_flags": {"person": True}, "court_polygon": [[0.1, 0.1], [0.9, 0.1],
[0.5, 0.9]],
47                     }}, fh)
48             cfg = load_config(p)
49             f = cfg.indexing.fusion
50             assert f.substrate == "tracknet"
51             assert f.min_crossings == 2
52             assert f.cue_flags == {"person": True}
53             assert len(f.court_polygon) == 3
```

```

54         # An old config without a fusion block still loads (defaults applied).
55         assert cfg.indexing.wasb.velocity_threshold == 0.02
56
57
58     def test_file_override():
59         with tempfile.TemporaryDirectory() as td:
60             p = os.path.join(td, "override.json")
61             [external_agent_tool_result]

```

49. assistant (2026-07-02T06:11:33.343Z)

```

[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\config\loader.py
[/external_agent_tool_call]

```

50. assistant (2026-07-02T06:11:33.343Z)

```

[external_agent_tool_result]
1     """Configuration loader and saver.
2
3     Provides a single place to load validated AppConfig with built-in defaults
4     merged with optional file overrides (config.json).
5
6     This replaces scattered json.load + .get() patterns throughout the codebase.
7     """
8
9     from __future__ import annotations
10
11     import json
12     import logging
13     from pathlib import Path
14     from typing import Any
15
16     from pydantic import BaseModel
17
18     from .models import AppConfig
19
20     logger = logging.getLogger(__name__)
21
22     DEFAULT_CONFIG_PATH = Path("config.json")
23
24
25     def _unknown_key_paths(data: dict[str, Any], model_cls: type[BaseModel], prefix: str =
26     "") -> list[str]:
27         """Dotted paths in `data` that the typed config will not honor.
28
29         Two silent failure modes exist without this check: a typo'd TOP-LEVEL key is
30         kept by AppConfig's ``extra="allow"`` but never read by typed code, and a
31         typo'd key INSIDE a section is silently dropped by Pydantic. Either way the
32         user's intent is lost ? collect every such key so load_config can warn.
33         """
34         unknown: list[str] = []
35         fields = model_cls.model_fields
36         for key, val in data.items():
37             if key not in fields:
38                 unknown.append(prefix + key)
39                 continue
40             ann = fields[key].annotation
41             if isinstance(val, dict) and isinstance(ann, type) and issubclass(ann,
42             BaseModel):
43                 unknown.extend(_unknown_key_paths(val, ann, f"{prefix}{key}.")
44         return unknown
45
46     def _get_builtin_defaults() -> dict[str, Any]:

```

```

46     """Return a plain dict of the current built-in defaults.
47
48     The AppConfig model is the single source of truth for defaults.
49     """
50     return AppConfig().model_dump(mode="json")
51
52
53 def load_config(path: str | Path | None = None) -> AppConfig:
54     """Load configuration with the following precedence:
55
56     1. Built-in defaults defined in the Pydantic models.
57     2. Values from the JSON file (if present and readable).
58     3. (Future) environment / CLI overrides.
59
60     Missing file or unreadable file ? returns a fully defaulted AppConfig.
61     Unknown keys in the file are tolerated (extra="allow" on the model).
62     """
63     cfg_path = Path(path) if path is not None else DEFAULT_CONFIG_PATH
64
65     file_data: dict[str, Any] = {}
66     if cfg_path.exists():
67         try:
68             with open(cfg_path, "r", encoding="utf-8") as f:
69                 file_data = json.load(f)
70         except Exception as exc:
71             # Non-fatal: continue with defaults + whatever we could parse.
72             # In production we might log here.
73             logger.warning(f"Failed to read {cfg_path}: {exc}. Using defaults + partial
data.")
74
75     if file_data:
76         unknown = _unknown_key_paths(file_data, AppConfig)
77         if unknown:
78             logger.warning(
79                 "[CONFIG WARNING] %s contains keys the typed config does not "
80                 "recognize (typo'd or removed knobs ? top-level extras are kept "
81                 "but unread; unknown section keys are silently dropped): %s",
82                 cfg_path, ", ".join(sorted(unknown)),
83             )
84
85     # Merge: start from defaults, overlay file data (Pydantic will validate/coerce).
86     merged = {**_get_builtin_defaults(), **file_data}
87
88     return AppConfig.model_validate(merged)
89
90
91 def save_default_config(path: str | Path | None = None) -> Path:
92     """Write a fresh default config.json.
93
94     Used by setup.py diagnostics and the --setup CLI path.
95     Overwrites if the file exists.
96     """
97     cfg_path = Path(path) if path is not None else DEFAULT_CONFIG_PATH
98     cfg = AppConfig()
99     cfg_path.parent.mkdir(parents=True, exist_ok=True)
100     cfg_path.write_text(cfg.model_dump_json(indent=4), encoding="utf-8")
101     return cfg_path
102
103
104     # Backwards-compatible alias u...
[/external_agent_tool_result]

```

51. assistant (2026-07-02T06:11:33.343Z)

[external_agent_tool_result]

```
1      # Utils package
2
[/external_agent_tool_result]
```

52. assistant (2026-07-02T06:11:33.343Z)

I have a complete picture. Let me verify `backend/\_\_init\_\_.py` is lightweight (so importing the logging helper early is safe) and confirm there's no existing `test\_export.py`.

53. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\__init__.py
[/external_agent_tool_call]
```

54. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_result: error]
File does not exist. Note: your current working directory is C:\Users\avidu\Projects\badminton-highlight-indexer.
[/external_agent_tool_result]
```

55. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\utils\__init__.py
[/external_agent_tool_call]
```

56. assistant (2026-07-02T06:11:33.343Z)

Branch is safe to set up. Let me create it from the **remote** master (not the current working branch).

57. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_call: Bash]
description: Create branch from origin/master
command: git checkout -b feat/runlog-readability origin/master && git branch --show-current &&
git log --oneline -1
[/external_agent_tool_call]
```

58. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_result]
branch 'feat/runlog-readability' set up to track 'origin/master'.
Switched to a new branch 'feat/runlog-readability'
feat/runlog-readability
35a50f3 docs(readme): index QUALITY_ITERATIONS + mark EXPERIMENT_HARNESS P1 SHIPPED (#118)
[/external_agent_tool_result]
```

59. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_result]
File created successfully at: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\utils\logging_setup.py (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

60. assistant (2026-07-02T06:11:33.343Z)

Branch is correctly based on `origin/master` @ 35a50f3. Here's my plan before I implement:

Plan

****Change 1 ? quieter `run.log` (HTTP chatter ? WARNING, config-gated)****

- New central helper `backend/utils/logging\_setup.py::setup\_logging(verbose\_http=False)` ? runs `basicConfig` (same format as today) and raises `httpx, httpcore, urllib3, google\_genai, google.genai, google.auth, grpc` to `WARNING` (NOTSET when `verbose\_http=True` to restore for debugging). Idempotent; never touches `backend.\*`.
- New config knob `logging.verbose\_http: bool = False` (`LoggingConfig` on `AppConfig`).
- Wire into the manual-QA run path `run\_process\_and\_stitch.py`: call `setup\_logging()` at import (quiet default), then `setup\_logging(verbose\_http=cfg.logging.verbose\_http)` right after config load.

****Change 2 ? end-of-run rally-times summary****

- New pure helper `backend/tools/export.py::format\_rally\_summary(segments) -> str` (sits next to `segments\_to\_csv`/`basic\_stats`, same play_segment dicts ? text): header (total rallies + total play in HH:MM:SS) + per-rally table `# start end dur(s) shot\_count`, HH:MM:SS via a local `\_fmt\_hms` matching the segmenters' `\_fmt\_ts`. Handles `metadata` as dict or JSON string.
- Log it as one block at the END of `run\_process\_and\_stitch.py::main()` so the log lines up row-for-row with the rally CSV.

****Tests:**** `tests/test\_logging\_setup.py`, `tests/test\_rally\_summary.py`, + a default/override assertion in `tests/test\_config.py`. Verify `python run\_all\_tests.py` green, ruff + mypy clean.

Implementing now.

61. assistant (2026-07-02T06:11:33.343Z)

```
[external_agent_tool_call: Write]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\utils\logging_setup.py
[/external_agent_tool_call]
```

62. assistant (2026-07-02T06:11:33.343Z)

Now the config knob:

63. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_call: Edit]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\config\models.py
[/external_agent_tool_call]
```

64. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_result]
The file C:\Users\avidu\Projects\badminton-highlight-indexer\backend\config\models.py has been
updated successfully. (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

65. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_call: Edit]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\config\models.py
[/external_agent_tool_call]
```

66. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_result]
The file C:\Users\avidu\Projects\badminton-highlight-indexer\backend\config\models.py has been
updated successfully. (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

67. assistant (2026-07-02T06:11:33.344Z)

Now the rally-summary helper in `export.py`:

68. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_call: Edit]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\tools\export.py
[/external_agent_tool_call]
```

69. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_result]
The file C:\Users\avidu\Projects\badminton-highlight-indexer\backend\tools\export.py has been
updated successfully. (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

70. assistant (2026-07-02T06:11:33.344Z)

Now wire both into the run path. First the logging setup at import + after config load:

71. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_call: Edit]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\run_process_and_stitch.py
[/external_agent_tool_call]
```

72. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_result]
The file C:\Users\avidu\Projects\badminton-highlight-indexer\run_process_and_stitch.py has been
updated successfully. (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

73. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_call: Edit]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\run_process_and_stitch.py
[/external_agent_tool_call]
```

74. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_result]
The file C:\Users\avidu\Projects\badminton-highlight-indexer\run_process_and_stitch.py has been
updated successfully. (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

75. assistant (2026-07-02T06:11:33.344Z)

Now add the summary block at the end of the run. Let me view the current end of `main()`:

76. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\run_process_and_stitch.py
[/external_agent_tool_call]
```

77. assistant (2026-07-02T06:11:33.344Z)

```
[external_agent_tool_result]
128         if skipped_count > 0:
129             print(f"[FILTER] Excluded {skipped_count} play segments with fewer than
{min_shot_count} shots.", flush=True)
130             print(f"[FILTER] Retaining {len(time_pairs)} play segments for stitching.",
flush=True)
131
132         if not time_pairs:
```

```

133         print("No play segments passed the filtering criteria. Cannot compile
highlights.")
134         return
135
136         compiler = VideoCompiler(output_dir, end_padding=end_padding,
begin_padding=begin_padding,
137                                   watermark=cfg.stitching.watermark)
138         try:
139             out_path = compiler.compile_highlights(video_path, time_pairs, output_filename)
140             print(f"Stitched highlight video saved to: {out_path}", flush=True)
141         except Exception as e:
142             print(f"Error during stitching: {e}", flush=True)
143
144     if __name__ == "__main__":
145         main()
146
[/external_agent_tool_result]

```

78. assistant (2026-07-02T06:11:33.344Z)

I'll emit the summary as the final block (after stitching, so ffmpeg output doesn't bury it), and also on the all-filtered early-return path so a fully-filtered run still shows what was produced:

79. assistant (2026-07-02T06:11:33.344Z)

```

[external_agent_tool_call: Edit]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\run_process_and_stitch.py
[/external_agent_tool_call]

```

80. assistant (2026-07-02T06:11:33.344Z)

```

[external_agent_tool_result]
The file C:\Users\avidu\Projects\badminton-highlight-indexer\run_process_and_stitch.py has been
updated successfully. (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]

```

81. assistant (2026-07-02T06:11:33.344Z)

Now the tests. First the rally-summary helper test:

82. assistant (2026-07-02T06:11:33.344Z)

```

[external_agent_tool_call: Write]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\tests\test_rally_summary.py
[/external_agent_tool_call]

```

83. assistant (2026-07-02T06:11:33.345Z)

```

[external_agent_tool_result]
File created successfully at: C:\Users\avidu\Projects\badminton-highlight-
indexer\tests\test_rally_summary.py (file state is current in your context ? no need to Read it
back)
[/external_agent_tool_result]

```

84. assistant (2026-07-02T06:11:33.345Z)

```

[external_agent_tool_call: Write]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\tests\test_logging_setup.py
[/external_agent_tool_call]

```

85. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_result]
File created successfully at: C:\Users\avidu\Projects\badminton-highlight-indexer\tests\test_logging_setup.py (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

86. assistant (2026-07-02T06:11:33.345Z)

Now add config assertions for the new knob to `test\_config.py`:

87. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_call: Edit]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\tests\test_config.py
[/external_agent_tool_call]
```

88. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_result]
The file C:\Users\avidu\Projects\badminton-highlight-indexer\tests\test_config.py has been updated successfully. (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

89. assistant (2026-07-02T06:11:33.345Z)

Now let me run the new/affected tests and the linters.

90. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_call: Bash]
description: Run new + config tests
command: python -m pytest tests/test_rally_summary.py tests/test_logging_setup.py
tests/test_config.py -q 2>&1 | tail -30
[/external_agent_tool_call]
```

91. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_result]
.....F..... [100%]
===== FAILURES =====
_____ test_does_not_touch_backend_loggers _____

restore_noisy_levels = None

    def test_does_not_touch_backend_loggers(restore_noisy_levels):
        # Our own loggers must keep emitting at INFO ? we only ever tune the named third
        parties.
        setup_logging()
        assert logging.getLogger("backend.pipeline.segmenters.trajectory_hybrid").level ==
logging.NOTSET
> assert logging.getLogger().isEnabledFor(logging.INFO)
E       assert False
E       + where False = isEnabledFor(20)
E       +   where isEnabledFor = <RootLogger root (WARNING)>.isEnabledFor
E       +     where <RootLogger root (WARNING)> = <function getLogger at 0x00000272EDEC1A80>()
E       +       where <function getLogger at 0x00000272EDEC1A80> = logging.getLogger
E       +         and 20 = logging.INFO

tests\test_logging_setup.py:42: AssertionError
===== short test summary info =====
FAILED tests/test_logging_setup.py::test_does_not_touch_backend_loggers - ass...
1 failed, 20 passed in 0.61s
[/external_agent_tool_result]
```


92. assistant (2026-07-02T06:11:33.345Z)

The root-level assertion fails only under pytest (pytest pre-installs a root handler, so `basicConfig` becomes a no-op for level). In the real entrypoint `basicConfig` runs first and does set INFO ? but the `level` param should take effect deterministically regardless. I'll force the root level explicitly (more robust) and make the test restore it.

93. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_call: Edit]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\backend\utils\logging_setup.py
[/external_agent_tool_call]
```

94. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_result]
The file C:\Users\avidu\Projects\badminton-highlight-indexer\backend\utils\logging_setup.py has
been updated successfully. (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

95. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_call: Edit]
file: C:\Users\avidu\Projects\badminton-highlight-indexer\tests\test_logging_setup.py
[/external_agent_tool_call]
```

96. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_result]
The file C:\Users\avidu\Projects\badminton-highlight-indexer\tests\test_logging_setup.py has
been updated successfully. (file state is current in your context ? no need to Read it back)
[/external_agent_tool_result]
```

97. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_call: Bash]
description: Re-run new + config tests
command: python -m pytest tests/test_rally_summary.py tests/test_logging_setup.py
tests/test_config.py -q 2>&1 | tail -15
[/external_agent_tool_call]
```

98. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_result]
..... [100%]
21 passed in 0.35s
[/external_agent_tool_result]
```

99. assistant (2026-07-02T06:11:33.345Z)

All green. Now ruff + mypy on the touched files, then the full suite.

100. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_call: Bash]
description: Ruff check on touched files
command: python -m ruff check backend/utils/logging_setup.py backend/config/models.py
backend/tools/export.py run_process_and_stitch.py tests/test_rally_summary.py
tests/test_logging_setup.py tests/test_config.py 2>&1 | tail -30
[/external_agent_tool_call]
```

101. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_result]
All checks passed!
[/external_agent_tool_result]
```

102. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_call: Bash]
description: Mypy on touched backend modules
command: python -m mypy backend/utils/logging_setup.py backend/tools/export.py
backend/config/models.py 2>&1 | tail -30
[/external_agent_tool_call]
```

103. assistant (2026-07-02T06:11:33.345Z)

```
[external_agent_tool_result]
Success: no issues found in 3 source files
[/external_agent_tool_result]
```

104. assistant (2026-07-02T06:11:33.346Z)

Let me check how CI runs mypy/ruff (to match scope) and run the full suite.

105. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_call: Bash]
description: Locate lint/type configs and CI invocation
command: git ls-files | grep -iE
"pyproject.toml|mypy.ini|setup.cfg|ruff.toml|.github/workflows" 2>&1 | head; echo "---- CI
mypy/ruff invocation ----"; grep -riE "mypy|ruff" .github/ pyproject.toml setup.cfg 2>/dev/null
| grep -v "^Binary" | head -30
[/external_agent_tool_call]
```

106. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_result]
.github/workflows/ci.yml
mypy.ini
ruff.toml
---- CI mypy/ruff invocation ----
.github/workflows/ci.yml: # Lint gate (ruff.toml at repo root): pyflakes/pycodestyle + bugbear.
.github/workflows/ci.yml: # Also the repo-wide syntax guard ? ruff reports invalid-syntax on
every
.github/workflows/ci.yml:     name: Ruff lint
.github/workflows/ci.yml:     - name: Install ruff
.github/workflows/ci.yml:       run: python -m pip install ruff
.github/workflows/ci.yml:     - name: Ruff check
.github/workflows/ci.yml:       run: ruff check . --output-format github
.github/workflows/ci.yml: # Type gate (mypy.ini at repo root): lenient green baseline with an
explicit
.github/workflows/ci.yml: # up. Typed core deps are installed so mypy sees their real
signatures
.github/workflows/ci.yml:     name: Mypy (lenient baseline)
.github/workflows/ci.yml:     - name: Install mypy + typed core deps
.github/workflows/ci.yml:       run: python -m pip install mypy pydantic fastapi uvicorn
python-dotenv google-genai numpy
.github/workflows/ci.yml:     - name: Mypy
.github/workflows/ci.yml:       run: mypy backend training
[/external_agent_tool_result]
```

107. assistant (2026-07-02T06:11:33.346Z)

CI runs `ruff check .` (whole repo) and `mypy backend training`. Let me match those exactly, plus the full test suite.

108. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_call: Bash]
description: Repo-wide ruff + CI-scope mypy
command: python -m ruff check . --output-format concise 2>&1 | tail -15; echo "==== MYPY =====";
python -m mypy backend training 2>&1 | tail -15
[/external_agent_tool_call]
```

109. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_result]
All checks passed!
==== MYPY =====
backend\annotations\base.py:85: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\sports\_init_.py:7: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\pipeline\validators\base.py:36: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\_init_.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 81 source files
[/external_agent_tool_result]
```

110. assistant (2026-07-02T06:11:33.346Z)

Ruff clean, mypy clean (the notes are pre-existing baseline noise, "no issues found"). Now the full suite:

111. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_call: Bash]
description: Run the full test suite
command: python run_all_tests.py 2>&1 | tail -25
[/external_agent_tool_call]
```

112. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_result]
=====
Full test suite: C:\Users\avidu\AppData\Local\Programs\Python\Python313\python.exe -m pytest
-q
obsreport (optional reporting dep): available
=====
..... [ 15%]
..... [ 30%]
..... [ 45%]
..... [ 60%]
..... [ 75%]
..... [ 90%]
..... [100%]
477 passed in 36.35s
=====
RESULT: PASS - all tests green
=====
[/external_agent_tool_result]
```

113. assistant (2026-07-02T06:11:33.346Z)

All 477 pass. Let me do a quick functional sanity check of the summary output and confirm the entrypoint compiles:

114. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_call: Bash]
description: Demo summary output + compile entrypoint
command: python -c "from backend.tools.export import format_rally_summary; print(format_rally_summary({'start_time':10.0,'end_time':25.4,'duration':15.4,'metadata':{'shot_count':8}},{'start_time':95.0,'end_time':3680.0,'duration':3585.0,'metadata':{'\shot_count\':11}}),{'start_time':3700.0,'end_time':3705.0,'duration':5.0}]))"; echo "---- compile entrypoint ----"; python -m py_compile run_process_and_stitch.py && echo "run_process_and_stitch.py compiles OK"
[/external_agent_tool_call]
```

115. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_result]
=== Rally summary: 3 rallies, 01:00:05 (3605.4s) total play ===
#  start      end      dur(s)  shots
1  00:00:10  00:00:25    15.4     8
2  00:01:35  01:01:20   3585.0    11
3  01:01:40  01:01:45     5.0     -
---- compile entrypoint ----
run_process_and_stitch.py compiles OK
[/external_agent_tool_result]
```

116. assistant (2026-07-02T06:11:33.346Z)

Output is clean and aligned, JSON-string metadata parsed (rally 2 ? 11 shots), missing shot_count ? `-.`. Now let me re-verify branch state and stage only my explicit paths.

117. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_call: Bash]
description: Verify branch and my changed files
command: git branch --show-current && git log --oneline -1 && echo "---- my changed files ----"
&& git status --short backend/utils/logging_setup.py backend/config/models.py
backend/tools/export.py run_process_and_stitch.py tests/test_rally_summary.py
tests/test_logging_setup.py tests/test_config.py
[/external_agent_tool_call]
```

118. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_result]
feat/runlog-readability
35a50f3 docs(readme): index QUALITY_ITERATIONS + mark EXPERIMENT_HARNESS P1 SHIPPED (#118)
---- my changed files ----
M backend/config/models.py
M backend/tools/export.py
M run_process_and_stitch.py
M tests/test_config.py
?? backend/utils/logging_setup.py
?? tests/test_logging_setup.py
?? tests/test_rally_summary.py
[/external_agent_tool_result]
```

119. assistant (2026-07-02T06:11:33.346Z)

```
[external_agent_tool_call: Bash]
description: Stage explicit paths and show what's staged
command: git add backend/utils/logging_setup.py backend/config/models.py backend/tools/export.py
run_process_and_stitch.py tests/test_rally_summary.py tests/test_logging_setup.py
tests/test_config.py && echo "---- staged ----" && git status --short && echo "---- diffstat ----"
&& git diff --cached --stat
[/external_agent_tool_call]
```

120. assistant (2026-07-02T06:11:33.346Z)

```

[external_agent_tool_result]
warning: in the working copy of 'backend/utils/logging_setup.py', LF will be replaced by CRLF
the next time Git touches it
warning: in the working copy of 'tests/test_logging_setup.py', LF will be replaced by CRLF the
next time Git touches it
warning: in the working copy of 'tests/test_rally_summary.py', LF will be replaced by CRLF the
next time Git touches it
---- staged ----
M backend/config/models.py
M backend/tools/export.py
A backend/utils/logging_setup.py
M run_process_and_stitch.py
M tests/test_config.py
A tests/test_logging_setup.py
A tests/test_rally_summary.py
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
?? scratch/arial.ttf
?? scratch/debug_wsl.py
?? scratch/keep_awake.ps1
?? scratch/probe_cv2_count.py
?? scratch/probe_gemini.py
?? scratch/repro_drawtext_font.py
?? scratch/repro_movie_overlay.py
?? scratch/run_gx010125.py
?? scratch/stitch_min5.py
?? scratch/trim_video.py
---- diffstat ----
 backend/config/models.py      | 10 ++++++
 backend/tools/export.py       | 59 ++++++
 backend/utils/logging_setup.py | 56 ++++++
 run_process_and_stitch.py     | 25 ++++++
 tests/test_config.py          | 15 ++++++
 tests/test_logging_setup.py   | 44 ++++++
 tests/test_rally_summary.py   | 56 ++++++
 7 files changed, 259 insertions(+), 6 deletions(-)
[/external_agent_tool_result]

```

121. assistant (2026-07-02T06:11:33.347Z)

Only my 7 files are staged; the sibling session's `docs/`, `artifacts/`, and `scratch/` changes are untouched. Committing now.

122. assistant (2026-07-02T06:11:33.347Z)

```

[external_agent_tool_call: Bash]
description: Commit the change
command: git commit -m '@'
feat(logging): quieter run.log ? suppress HTTP chatter + end-of-run rally summary

```

Two manual-QA readability changes for run.log (the owner runs the tool on a video, then eyeballs the rally CSV against the log to verify quality).

- 1) Suppress third-party HTTP/RPC chatter. The Gemini path (google-genai -> httpx/grpc/google.auth) logs one INFO line per request, burying our own [rally]/[compile]/[CACHE] lines. New central helper backend/utils/logging_setup.py::setup_logging raises httpx, httpcore, urllib3, google_genai/google.genai, google.auth, grpc to WARNING by default, gated behind a new config knob logging.verbose_http (false). Our own backend.* loggers are untouched. Wired into run_process_and_stitch.py.

2) End-of-run rally-times summary. New pure helper
backend/tools/export.py::format_rally_summary(segments) -> str renders a
header (total rallies + total play in HH:MM:SS) plus a per-rally table
(# start end dur(s) shot_count) from the produced play_segments, so the log
lines up row-for-row with the rally CSV. Logged as one block at the end of
run_process_and_stitch.py::main(). Timestamps use HH:MM:SS to match _fmt_ts.

Tests: tests/test_logging_setup.py, tests/test_rally_summary.py, +
logging.verbose_http default/override in tests/test_config.py.
Full suite green (477), ruff + mypy clean. Default behavior = quieter log +
the summary block.

Co-Authored-By: Claude Opus 4.8 (1M context) <noreply@anthropic.com>
'@
[/external_agent_tool_call]

123. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_result]
[feat/runlog-readability 80643b8] @ feat(logging): quieter run.log ? suppress HTTP chatter +
end-of-run rally summary
7 files changed, 259 insertions(+), 6 deletions(-)
create mode 100644 backend/utls/logging_setup.py
create mode 100644 tests/test_logging_setup.py
create mode 100644 tests/test_rally_summary.py
[/external_agent_tool_result]

124. assistant (2026-07-02T06:11:33.347Z)

The commit message got mangled ? I used PowerShell here-string syntax (`@'...'@`) inside the
Bash tool, which treats `@` literally. Let me inspect and fix it.

125. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_call: Bash]
description: Show raw commit message start
command: git log -1 --format='%B' | head -5 | cat -A | head -5
[/external_agent_tool_call]

126. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_result]
@\$
feat(logging): quieter run.log M-bM-^@M-^T suppress HTTP chatter + end-of-run rally summary\$
\$
Two manual-QA readability changes for run.log (the owner runs the tool on a\$
video, then eyeballs the rally CSV against the log to verify quality).\$
[/external_agent_tool_result]

127. assistant (2026-07-02T06:11:33.347Z)

Confirmed ? leading `@` line and trailing `@`. I'll fix it cleanly by writing the message to a
temp file and amending with `-F`.

128. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_call: Write]
file: C:\Users\avidu\AppData\Local\Temp\runlog_commit_msg.txt
[/external_agent_tool_call]

129. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_result]
File created successfully at: C:\Users\avidu\AppData\Local\Temp\runlog_commit_msg.txt (file

state is current in your context ? no need to Read it back)
[/external_agent_tool_result]

130. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_call: Bash]
description: Amend commit with clean message from file
command: git commit --amend -F "C:/Users/avidu/AppData/Local/Temp/runlog_commit_msg.txt" && echo
"---- verify ----" && git log -1 --format='%s' && git log --oneline origin/master..HEAD
[/external_agent_tool_call]

131. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_result]
[feat/runlog-readability 3af7d1a] feat(logging): quieter run.log ? suppress HTTP chatter + end-
of-run rally summary
Date: Sun Jun 14 01:28:03 2026 +0530
7 files changed, 259 insertions(+), 6 deletions(-)
create mode 100644 backend/utils/logging_setup.py
create mode 100644 tests/test_logging_setup.py
create mode 100644 tests/test_rally_summary.py
---- verify ----
feat(logging): quieter run.log ? suppress HTTP chatter + end-of-run rally summary
3af7d1a feat(logging): quieter run.log ? suppress HTTP chatter + end-of-run rally summary
[/external_agent_tool_result]

132. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_result]
File created successfully at: C:\Users\avidu\AppData\Local\Temp\runlog_pr_body.md (file state is
current in your context ? no need to Read it back)
[/external_agent_tool_result]

133. assistant (2026-07-02T06:11:33.347Z)

Clean message, and `origin/master..HEAD` shows only my single commit. Now push and open the PR
(the deliverable calls for a plan in the PR description). Let me write the PR body to a file to
avoid quoting pitfalls.

134. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_call: Write]
file: C:\Users\avidu\AppData\Local\Temp\runlog_pr_body.md
[/external_agent_tool_call]

135. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_call: Bash]
description: Push branch to origin
command: git push -u origin feat/runlog-readability 2>&1 | tail -8
[/external_agent_tool_call]

136. assistant (2026-07-02T06:11:33.347Z)

[external_agent_tool_result]
remote:
remote: Create a pull request for 'feat/runlog-readability' on GitHub by visiting:
remote: <https://github.com/avidullu/badminton-highlight-indexer/pull/new/feat/runlog-readability>
remote:
branch 'feat/runlog-readability' set up to track 'origin/feat/runlog-readability'.
To <https://github.com/avidullu/badminton-highlight-indexer.git>
* [new branch] feat/runlog-readability -> feat/runlog-readability
[/external_agent_tool_result]

137. assistant (2026-07-02T06:11:33.347Z)

```
[external_agent_tool_call: Bash]
description: Open the PR
command: gh pr create --base master --head feat/runlog-readability --title "feat(logging):
quieter run.log ? suppress HTTP chatter + end-of-run rally summary" --body-file
"C:/Users/avidu/AppData/Local/Temp/runlog_pr_body.md" 2>&1 | tail -5
[/external_agent_tool_call]
```

138. assistant (2026-07-02T06:11:33.348Z)

```
[external_agent_tool_result]
https://github.com/avidullu/badminton-highlight-indexer/pull/120
[/external_agent_tool_result]
```

139. assistant (2026-07-02T06:11:33.348Z)

PR opened: **https://github.com/avidullu/badminton-highlight-indexer/pull/120**

Per the checkpoint convention, let me update the live current-state memory to record this PR.
First I'll read it to append accurately.

140. assistant (2026-07-02T06:11:33.348Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md
[/external_agent_tool_call]
```

141. assistant (2026-07-02T06:11:33.348Z)

```
[external_agent_tool_result]
<system-reminder>[Truncated: PARTIAL view ? showing lines 1-461 of 1667 total (76837 tokens, cap
25000). Call Read with offset=462 limit=461 for the next page, or Grep to find a specific
section. Do NOT answer from this page alone if the answer may be further in the file.]</system-
reminder>
```

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4      metadata:
5        node_type: memory
6        type: project
7        originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-14 (ANALYZERS DOC ? FRAMEWORK + paper framing ? **PR #119 OPEN,
owner-gated, do NOT merge**; branch
13     `docs/analyzers-framework` off master `35a50f3`; docs-only). Owner fed two more ideas
this session and asked to fold them in
14     and treat the doc as a **reusable framework / research paper**, not a project patch
(ties [[paper-coauthor-aim]]). Extended
15     the on-master `docs/ANALYZERS_AND_RECONCILERS.md` (which had landed via #117) from "two
seams" into a **multi-scope
16     signal-fusion framework**: producers emit SIGNALS at **3 temporal scopes** ? frame
detectors (`RallyCue`) / window analyzers
17     / **session analyzers** ? all consumed as feature columns by the learned scorer (spine
unchanged: **no special-casing**, signal
18     ?scorer, never `if/else`). Folded in: (1) **warm-up/practice** as a new
**`SessionAnalyzer`** tier (whole-timeline span ? an
19     `in_warmup` membership feature the scorer DOWN-WEIGHTS, NOT a hard `if t<warmup_end:
```



```

drop`; measurable on golden TODAY since
20     warm-up is already labeled negative); (2) **shuttle-state**
(`in_air`/`racket_contact`/`inactive`) placed correctly as a
21     **frame-scope detector** (a MULTI_SIGNAL_FUSION_PLAN `RallyCue`, NOT a window analyzer)
? derive-from-existing-signals first
22     (velocity+visibility+direction-reversals+audio thwack), dedicated per-frame model ONLY
where WASB trajectory gaps bite, with
23     the per-frame-label blocker flagged. Added **$10 generality & research positioning**
(weakly-supervised temporal segmentation;
24     standard-vs-contributed) with an explicit **attribution honesty caveat**: NO fabricated
citations/results, related-work +
25     ablations + >1-sport still owed, n=6 = direction-check only ? the honest backbone the
paper aim needs. Diff = `2-seam?framework`
26     (310+/137?). ? Shared checkout had flipped to master mid-task (the sibling's collision
lesson recurred) ? re-verified branch,
27     branched fresh, normal commit (no in-flight WIP left to isolate; #117/#118 already
merged so all cross-links resolve). PR #116
28     stays CLOSED (content on master via #117). ?? NEXT (owner-gated): review #119; impl
tracks unchanged (fusion P3 LOVO, then
29     analyzers/reconcilers A1?A5 / R1?R3 per the doc's phasing ? each ONE PR, default OFF,
promote on measured LOVO lift).
30
31     **Checkpoint:** 2026-06-14 (FUSION P2 SHIPPED + collision cleanup + paper aim ? master
`35a50f3`; suite 467 green; 0 open PRs).
32     **PR #117 MERGED = `FusionSegmenter`** (fusion P2, registry `fusion_hybrid`, **default-
OFF**): shuttle-primary (windows seed
33     from WASB/TrackNet) + adjudicating cues ? `net_crossings` (Gate A, always persisted,
GPU-free via rally_gate w/ a reused
34     axis estimate) + optional person cue (Gate B) persisting the ~5 `fusion_features.py`
features (pure-tested) via
35     `PersonDetector.detections_per_frame` over each window's ORIGINAL-video frame range.
Plugs in via **two IDENTITY hooks** on
36     `trajectory_hybrid` (`_enrich_candidate_stats`/`_select_for_handoff`) ? wasb/tracknet
byte-identical (adversarially verified).
37     Served gating optional+default-OFF (`fusion.min_crossings`/`min_players`) ? gates the
handoff while persisting the full
38     SUPERSET (offline substrate intact). Court mask = optional `fusion.court_polygon`.
**Design-first** (`fusion-p2-design`
39     workflow) ? **31-agen...
[/external_agent_tool_result]

```

142. assistant (2026-07-02T06:11:33.348Z)

This is the live checkpoint log (newest first). I'll prepend a concise entry for PR #120. Let me insert it right after the header.

143. assistant (2026-07-02T06:11:33.348Z)

```

[external_agent_tool_call: Edit]
file: C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md
[/external_agent_tool_call]

```

144. assistant (2026-07-02T06:11:33.348Z)

```

[external_agent_tool_result: error]
<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>
[/external_agent_tool_result]

```

145. assistant (2026-07-02T06:11:33.348Z)

```

[external_agent_tool_call: Read]
file: C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md

```

```
[/external_agent_tool_call]
```

146. assistant (2026-07-02T06:11:33.348Z)

```
[external_agent_tool_result]
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-14 (ANALYZERS DOC ? FRAMEWORK + paper framing ? **PR #119 OPEN,
owner-gated, do NOT merge**); branch
13     `docs/analyzers-framework` off master `35a50f3`; docs-only). Owner fed two more ideas
this session and asked to fold them in
[/external_agent_tool_result]
```

147. assistant (2026-07-02T06:11:33.348Z)

```
[external_agent_tool_call: Edit]
file: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md
[/external_agent_tool_call]
```

148. assistant (2026-07-02T06:11:33.348Z)

```
[external_agent_tool_result]
The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md has been updated successfully. (file state is current in your
context ? no need to Read it back)
[/external_agent_tool_result]
```

149. assistant (2026-07-02T06:11:33.348Z)

Done. Both changes are implemented, tested, and shipped as a single PR off `master`.

PR: [#120 ? quieter `run.log` + end-of-run rally summary](<https://github.com/avidullu/badminton-highlight-ind>)

Branch `feat/runlog-readability` off `origin/master` @ `35a50f3`; one commit; `git log origin/master..HEAD` shows only mine. Staged explicit paths only ? the sibling session's `docs/`, `artifacts/`, `scratch/` changes are untouched. **Not merged (owner-gated).**

****Change 1 ? suppress HTTP chatter (config-gated)****

- New central helper `[backend/utils/logging_setup.py](backend/utils/logging_setup.py)` `::setup\_logging(verbose\_http=False)` ? raises `httpx`, `httpcore`, `urllib3`, `google_genai/google.genai`, `google.auth`, `grpc` to `WARNING` (the per-request lines from the Gemini path); ``verbose_http=True` ? `NOTSET`` to restore for debugging. ``backend.*`` left alone.
- New knob ``logging.verbose_http: false`` (`[backend/config/models.py](backend/config/models.py)`).
- Wired into `[run_process_and_stitch.py](run_process_and_stitch.py)` at import (quiet default) and re-applied after config load.

****Change 2 ? end-of-run rally summary****

- New pure helper `[backend/tools/export.py](backend/tools/export.py)` `::format\_rally\_summary(segments) -> str`, next to ``segments_to_csv/`basic_stats``. Sourced from the produced `play_segments` so it lines up row-for-row with the rally CSV. Sample:

```
...
```

```
=== Rally summary: 3 rallies, 01:00:05 (3605.4s) total play ===
```

#	start	end	dur(s)	shots
1	00:00:10	00:00:25	15.4	8
2	00:01:35	01:01:20	3585.0	11
3	01:01:40	01:01:45	5.0	-

...

- Logged as one block at the end of `main()` (and on the all-filtered early-return). HH:MM:SS via a local `\_fmt\_hms` matching the segmenters' `\_fmt\_ts`.

****Verification:**** `python run\_all\_tests.py` ? ****477 passed****; `ruff check .` clean; `mypy backend training` clean. New tests: `test\_logging\_setup.py`, `test\_rally\_summary.py`, plus default/override in `test\_config.py`.

****Two notes:****

- I scoped the wiring to `run\_process\_and\_stitch.py` (the manual-QA "run the tool on a video" path) to stay minimal. The FastAPI server (`backend/main.py`) still has its own `basicConfig` and isn't quieted ? say the word if you want the helper wired there too.
- Heads-up for the record: my first commit-message attempt got mangled (I used PowerShell `@'?'@` here-string syntax inside the ****Bash**** tool, which takes `@` literally). I caught it and amended cleanly via `!F file`; the final message is correct.

Machine-side check is yours (the dev container has no GPU/torch/cv2): run it on a real video and confirm the log is quieter and the summary block reads cleanly against the CSV.

150. user (2026-07-02T06:11:33.348Z)

can you merge the PR as well and ensure the master is clean from your side. I want to resolve this branch collisio

151. assistant (2026-07-02T06:11:33.348Z)

The owner has asked me to merge #120, and this is exactly the kind of good-engineering-fix that [[working-agreement-merging]] authorizes me to self-merge on verified-green CI ? now with an explicit request on top.

Quick note on approach: ultracode is on, but a merge + working-tree cleanup on a ****shared checkout**** is a serial, stateful git operation. Fanning out parallel agents here would race on `.git/index`/`HEAD` and risk clobbering the concurrent session's uncommitted WIP ? the opposite of what's wanted. So I'll do this directly and carefully, with thorough read-only verification at each step. Let me first establish the exact state (read-only).

152. assistant (2026-07-02T06:11:33.348Z)

```
[external_agent_tool_call: Bash]
description: Inspect PR #120 CI/merge state and local tree
command: echo "==== PR #120 state ===="; gh pr view 120 --json
number,state,mergeable,mergeStateStatus,headRefName,baseRefName,isCrossRepository 2>&1; echo;
echo "==== PR #120 checks ===="; gh pr checks 120 2>&1; echo; echo "==== local branch ===="; git
branch --show-current; echo; echo "==== git status (porcelain) ===="; git status --porcelain=v1;
echo; echo "==== staged diff stat ===="; git diff --cached --stat
[/external_agent_tool_call]
```

153. assistant (2026-07-02T06:11:33.349Z)

```
[external_agent_tool_result]
==== PR #120 state ====
{"baseRefName":"master","headRefName":"feat/runlog-readability","isCrossRepository":false,"merge
StateStatus":"CLEAN","mergeable":"MERGEABLE","number":120,"state":"OPEN"}

==== PR #120 checks ====
Full test suite (offline, mocked)      pass    1m31s    https://github.com/avidullu/badminton-
highlight-indexer/actions/runs/27477491781/job/81219327487
Mypy (lenient baseline) pass      34s      https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27477491781/job/81219327498
Ruff lint      pass      9s      https://github.com/avidullu/badminton-highlight-
```

```
indexer/actions/runs/27477491781/job/81219327531
Syntax sweep + import smoke (no GPU stack)      pass    21s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27477491781/job/81219327484
```

```
==== local branch ====
feat/runlog-readability
```

```
==== git status (porcelain) ====
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
?? scratch/arial.ttf
?? scratch/debug_wsl.py
?? scratch/keep_awake.ps1
?? scratch/p3_smoke.py
?? scratch/probe_cv2_count.py
?? scratch/probe_gemini.py
?? scratch/repro_drawtext_font.py
?? scratch/repro_movie_overlay.py
?? scratch/run_gx010125.py
?? scratch/stitch_min5.py
?? scratch/trim_video.py
```

```
==== staged diff stat ====
[/external_agent_tool_result]
```

154. assistant (2026-07-02T06:11:33.349Z)

CI is fully green (all 4 checks pass), `mergeable: CLEAN`, and the working tree now has **no tracked-file modifications** ? only the intentional untracked carryovers (`scratch/`, `artifacts/`). The sibling's previously-staged `docs/ANALYZERS\_AND\_RECONCILERS.md` change is no longer in the tree, so there's no WIP of theirs for me to preserve in the index. Let me re-fetch and confirm `origin/master` hasn't moved under me before the merge.

155. assistant (2026-07-02T06:11:33.349Z)

```
[external_agent_tool_call: Bash]
description: Confirm origin/master unmoved and branch diverge state
command: git fetch origin 2>&1; echo "==== origin/master tip ===="; git log --oneline -1
origin/master; echo "==== my branch tip ===="; git log --oneline -1 feat/runlog-readability;
echo "==== commits I'm adding (origin/master..HEAD) ===="; git log --oneline
origin/master..HEAD; echo "==== is my branch a clean fast-forward onto origin/master?
(master..mybranch should be just my 1 commit; origin/master..nothing-behind) ===="; echo "behind
origin/master by:"; git rev-list --count HEAD..origin/master
[/external_agent_tool_call]
```

156. assistant (2026-07-02T06:11:33.349Z)

```
[external_agent_tool_result]
==== origin/master tip ====
35a50f3 docs(readme): index QUALITY_ITERATIONS + mark EXPERIMENT_HARNESS P1 SHIPPED (#118)
==== my branch tip ====
3af7d1a feat(logging): quieter run.log ? suppress HTTP chatter + end-of-run rally summary
==== commits I'm adding (origin/master..HEAD) ====
3af7d1a feat(logging): quieter run.log ? suppress HTTP chatter + end-of-run rally summary
==== is my branch a clean fast-forward onto origin/master? (master..mybranch should be just my 1
commit; origin/master..nothing-behind) ====
behind origin/master by:
0
[/external_agent_tool_result]
```

157. assistant (2026-07-02T06:11:33.349Z)

Clean linear state confirmed: `origin/master` unmoved at `35a50f3`, my branch is exactly `origin/master` + my single commit `3af7d1a`, 0 behind. CI green, `mergeable: CLEAN`. Squash-merging now (repo convention is squash), deleting the branch on both sides.

158. assistant (2026-07-02T06:11:33.349Z)

```
[external_agent_tool_call: Bash]
description: Squash-merge PR #120 and delete the branch
command: gh pr merge 120 --squash --delete-branch 2>&1
[/external_agent_tool_call]
```

159. assistant (2026-07-02T06:11:33.349Z)

```
[external_agent_tool_result]
From https://github.com/avidullu/badminton-highlight-indexer
* branch          master      -> FETCH_HEAD
   35a50f3..6076eef master    -> origin/master
Updating 35a50f3..6076eef
Fast-forward
 backend/config/models.py      | 10 ++++++
 backend/tools/export.py       | 59 ++++++++++++++++++++++++++++++++++++++
 backend/utils/logging_setup.py | 56 ++++++++++++++++++++++++++++++++++++++
 run_process_and_stitch.py     | 25 ++++++-----
 tests/test_config.py          | 15 ++++++++
 tests/test_logging_setup.py   | 44 ++++++++++++++++++++++++++++++++++
 tests/test_rally_summary.py   | 56 ++++++++++++++++++++++++++++++++++++++
 7 files changed, 259 insertions(+), 6 deletions(-)
 create mode 100644 backend/utils/logging_setup.py
 create mode 100644 tests/test_logging_setup.py
 create mode 100644 tests/test_rally_summary.py
[/external_agent_tool_result]
```

160. assistant (2026-07-02T06:11:33.349Z)

Merge succeeded ? master advanced `35a50f3..6076eef` and fast-forwarded locally. Let me verify the complete final state: branch, sync, branch deletion (both sides), PR status, and working tree.

161. assistant (2026-07-02T06:11:33.349Z)

```
[external_agent_tool_call: Bash]
description: Verify final state after merge
command: echo "==== current branch ===="; git branch --show-current; echo; echo "==== local
master vs origin/master ===="; git log --oneline -1 HEAD; git log --oneline -1 origin/master;
echo; echo "==== HEAD ahead/behind origin/master (should be 0 0) ===="; echo "ahead:"; git rev-
list --count origin/master..HEAD; echo "behind:"; git rev-list --count HEAD..origin/master;
echo; echo "==== local branches (feat/runlog-readability should be GONE) ===="; git branch;
echo; echo "==== remote branch gone? (grep should be empty) ===="; git branch -r | grep runlog
|| echo " (no remote runlog branch ? good)"; echo; echo "==== PR #120 state ===="; gh pr view
120 --json state,mergedAt,mergeCommit --jq '{state, mergedAt, mergeCommit: .mergeCommit.oid}'
2>&1; echo; echo "==== working tree (porcelain) ===="; git status --porcelain=v1
[/external_agent_tool_call]
```

162. assistant (2026-07-02T06:11:33.349Z)

```
[external_agent_tool_result]
==== current branch ====
master

==== local master vs origin/master ====
6076eef feat(logging): quieter run.log ? suppress HTTP chatter + end-of-run rally summary (#120)
6076eef feat(logging): quieter run.log ? suppress HTTP chatter + end-of-run rally summary (#120)
```

```

==== HEAD ahead/behind origin/master (should be 0 0) ====
ahead:
0
behind:
0

==== local branches (feat/runlog-readability should be GONE) ====
ci/ruff-lint-lane
docs/analyzers-and-reconcilers
docs/analyzers-framework
docs/pmf-field-signals
docs/storage-sharing-model
docs/video-locality-model
fix/stitching-padding-config
* master

==== remote branch gone? (grep should be empty) ====
origin/feat/runlog-readability

==== PR #120 state ====
{"mergeCommit":"6076eef1f6c70a6ad49af7100a704b7456b694c","mergedAt":"2026-06-13T20:02:39Z","state":"MERGED"}

==== working tree (porcelain) ====
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
?? scratch/arial.ttf
?? scratch/debug_wsl.py
?? scratch/keep_awake.ps1
?? scratch/p3_smoke.py
?? scratch/probe_cv2_count.py
?? scratch/probe_gemini.py
?? scratch/repro_drawtext_font.py
?? scratch/repro_movie_overlay.py
?? scratch/run_gx010125.py
?? scratch/stitch_min5.py
?? scratch/trim_video.py
[/external_agent_tool_result]

```

163. assistant (2026-07-02T06:11:33.349Z)

Core objective achieved: I'm on `master`, local `master == origin/master == 6076eef` (the `(#120)` squash commit), my local feature branch is deleted, PR #120 is **MERGED**, and the working tree has only the intentional untracked carryovers ? no tracked changes, no leftover of mine.

Two cleanup items remain: a stale `origin/feat/runlog-readability` remote-tracking ref (the remote branch was deleted by the merge, but my local ref needs pruning), and a pile of **other sessions'** stale local branches. Let me prune the stale ref and map every local branch to its PR state before touching anything ? I won't delete branches tied to open PRs.

164. assistant (2026-07-02T06:11:33.349Z)

```

[external_agent_tool_call: Bash]
description: Prune stale ref and map local branches to PR states
command: echo "==== does the remote branch still exist on GitHub? (expect 404) ====="; gh api
repos/avidullu/badminton-highlight-indexer/branches/feat/runlog-readability --jq '.name' 2>&1 |
head -3; echo; echo "==== prune stale remote-tracking refs ====="; git remote prune origin 2>&1;
echo; echo "==== remote runlog ref after prune (expect empty) ====="; git branch -r | grep runlog
|| echo " (gone ? clean)"; echo; echo "==== map each LOCAL branch -> PR state ====="; for b in

```

```
ci/ruff-lint-lane docs/analyzers-and-reconcilers docs/analyzers-framework docs/pmf-field-signals
docs/storage-sharing-model docs/video-locality-model fix/stitching-padding-config; do printf
"%-38s " "$b"; gh pr list --state all --head "$b" --json number,state,title --jq '[] | "PR
#(.number) \(.state): \(.title)'" 2>/dev/null | head -1 || echo "(no PR)"; done
[/external_agent_tool_call]
```

165. assistant (2026-07-02T06:11:33.349Z)

```
[external_agent_tool_result]
==== does the remote branch still exist on GitHub? (expect 404) ====
{"message":"Branch not
found","documentation_url":"https://docs.github.com/rest/branches/branches#get-a-
branch","status":"404"}gh: Branch not found (HTTP 404)

==== prune stale remote-tracking refs ====
Pruning origin
URL: https://github.com/avidullu/badminton-highlight-indexer.git
* [pruned] origin/docs/readme-index-sync
* [pruned] origin/feat/experiment-harness
* [pruned] origin/feat/person-cue-extraction
* [pruned] origin/feat/runlog-readability

==== remote runlog ref after prune (expect empty) ====
(gone ? clean)

==== map each LOCAL branch -> PR state ====
ci/ruff-lint-lane PR #92 MERGED: ci: ruff lint lane + fix all violations;
remove never-compiled flyer stub
docs/analyzers-and-reconcilers PR #116 CLOSED: docs: two rally-detection extensibility
seams (injectable analyzers + reconciliation pass)
docs/analyzers-framework PR #119 OPEN: docs: extend analyzers/reconcilers into a
multi-scope signal-fusion framework (warm-up + shuttle-state)
docs/pmf-field-signals PR #104 MERGED: docs(pmf): C13 field kit ? recruiting
flyer, visit playbook, answer sheet, pre-registered signals
docs/storage-sharing-model PR #111 MERGED: docs(strategy): storage-priced sharing
mental model (DRAFT, co-design)
docs/video-locality-model PR #106 MERGED: docs(platform): video-locality mental
model ? the 10GB question
fix/stitching-padding-config PR #102 MERGED: fix(api): thread config stitching
paddings into live VideoCompiler sites
[/external_agent_tool_result]
```

166. assistant (2026-07-02T06:11:33.349Z)

Clear mapping. The remote branch is confirmed deleted on GitHub (404), and I pruned 4 stale remote-tracking refs (including mine). Now the local branches:

```
| Local branch | PR | State | Action |
|---|---|---|---|
| `ci/ruff-lint-lane` | #92 | MERGED | delete |
| `docs/analyzers-and-reconcilers` | #116 | CLOSED (redundant; content on master via #117) |
delete |
| `docs/pmf-field-signals` | #104 | MERGED | delete |
| `docs/storage-sharing-model` | #111 | MERGED | delete |
| `docs/video-locality-model` | #106 | MERGED | delete |
| `fix/stitching-padding-config` | #102 | MERGED | delete |
| `docs/analyzers-framework` | **#119 OPEN** | **active, owner-gated** | **KEEP** |
```

This matches your established protocol (prune local branches that map to merged/closed PRs; keep `master` + open-PR branches) and directly resolves the collision clutter. Branch deletion is local-only and reflog-recoverable. Deleting the 6 stale ones, keeping `master` + the live #119 branch.

167. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_call: Bash]
description: Delete stale merged/closed local branches, keep open-PR branch
command: echo "==== safety: confirm on master ===="; git branch --show-current; echo; echo "====
delete the 6 stale local branches (merged/closed PRs) ====="; git branch -D ci/ruff-lint-lane
docs/analyzers-and-reconcilers docs/pmf-field-signals docs/storage-sharing-model docs/video-
locality-model fix/stitching-padding-config 2>&1; echo; echo "===== remaining local branches
(expect: docs/analyzers-framework + master) ====="; git branch
[/external_agent_tool_call]
```

168. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_result]
==== safety: confirm on master ====
master

==== delete the 6 stale local branches (merged/closed PRs) =====
Deleted branch ci/ruff-lint-lane (was 75659d1).
Deleted branch docs/analyzers-and-reconcilers (was 0f673ff).
Deleted branch docs/pmf-field-signals (was 9c84fea).
Deleted branch docs/storage-sharing-model (was adbe579).
Deleted branch docs/video-locality-model (was 85abe26).
Deleted branch fix/stitching-padding-config (was 30e4899).

===== remaining local branches (expect: docs/analyzers-framework + master) =====
docs/analyzers-framework
* master
[/external_agent_tool_result]
```

169. assistant (2026-07-02T06:11:33.350Z)

All 6 stale branches deleted; only `master` (current) and `docs/analyzers-framework` (#119 OPEN) remain locally. Final verification of the whole clean state:

170. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_call: Bash]
description: Final clean-state verification
command: echo "==== on branch ====="; git branch --show-current; echo; echo "===== master tip
(should be #120 squash) + last 3 ====="; git log --oneline -3; echo; echo "===== sync with
origin/master (0 ahead / 0 behind) ====="; echo "ahead:"; git rev-list --count
origin/master..HEAD; echo "behind:"; git rev-list --count HEAD..origin/master; echo; echo "=====
local branches ====="; git branch; echo "===== remote branches still on origin ====="; git branch
-r; echo; echo "===== working tree (only untracked carryovers expected) ====="; git status
--short; echo; echo "===== merged code present on master? ====="; test -f
backend/utils/logging_setup.py && echo " backend/utils/logging_setup.py ? on master"; git show
HEAD:backend/tools/export.py | grep -c "def format_rally_summary" | xargs -I{} echo "
format_rally_summary defs on master HEAD: {}"
[/external_agent_tool_call]
```

171. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_result]
==== on branch =====
master

===== master tip (should be #120 squash) + last 3 =====
6076eef feat(logging): quieter run.log ? suppress HTTP chatter + end-of-run rally summary (#120)
35a50f3 docs(readme): index QUALITY_ITERATIONS + mark EXPERIMENT_HARNESS P1 SHIPPED (#118)
cbf2236 feat(segmenters): FusionSegmenter ? shuttle + net-crossing + person-cue fusion (P2,
default-OFF) (#117)

===== sync with origin/master (0 ahead / 0 behind) =====
ahead:
0
```



```

behind:
0

==== local branches ====
docs/analyzers-framework
* master
==== remote branches still on origin ====
origin/HEAD -> origin/master
origin/ci/ruff-lint-lane
origin/claude/docs-handoff-consistency
origin/claude/highlight-watermark
origin/claude/indexer-doc-review-XVPTh
origin/claude/platform-architecture-research
origin/claude/watermark-bundled-font
origin/docs/analyzers-and-reconcilers
origin/docs/analyzers-framework
origin/docs/next-steps-multisport
origin/docs/pmf-field-signals
origin/docs/storage-sharing-model
origin/docs/video-locality-model
origin/feat/async-job-model
origin/feat/fusion-segmenter
origin/feat/khelacharya-shot-intelligence
origin/feat/owned-model-study-and-prototype
origin/fix/stitching-padding-config
origin/master

==== working tree (only untracked carryovers expected) ====
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
?? scratch/arial.ttf
?? scratch/debug_wsl.py
?? scratch/keep_awake.ps1
?? scratch/p3_smoke.py
?? scratch/probe_cv2_count.py
?? scratch/probe_gemini.py
?? scratch/repro_drawtext_font.py
?? scratch/repro_movie_overlay.py
?? scratch/run_gx010125.py
?? scratch/stitch_min5.py
?? scratch/trim_video.py

==== merged code present on master? ====
backend/utils/logging_setup.py ? on master
format_rally_summary defs on master HEAD: 1
[/external_agent_tool_result]

```

172. assistant (2026-07-02T06:11:33.350Z)

My side is fully clean: on `master` at the `(#120)` squash commit, 0 ahead/0 behind `origin/master`, only `master` + the live `docs/analyzers-framework` (#119) remain locally, working tree has nothing but the intentional untracked carryovers, and the merged code is confirmed on master.

One thing surfaced that's **your call**, not mine to assume: the **remote** (`origin`) still carries 16 stale branches. Most map to already-merged/closed PRs, but a few are deliberately kept or still active. Deleting remote branches is outward-facing on the shared repo, so let me map them precisely before touching anything.

173. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_call: Bash]
description: Map all remote branches to PR states
command: echo "branch | PR/state"; for b in $(git branch -r | sed 's/ *origin\\/' | grep -vE
'^HEAD|^master'); do printf "%-44s " "$b"; pr=$(gh pr list --state all --head "$b" --json
number,state --jq '.[0] | "PR #\(.number) \(.state)"' 2>/dev/null); [ -n "$pr" ] && echo "$pr"
|| echo "(no PR ? branch only)"; done
[/external_agent_tool_call]
```

174. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_result]
branch | PR/state
ci/ruff-lint-lane                PR #92 MERGED
claude/docs-handoff-consistency  PR #59 MERGED
claude/highlight-watermark      PR #101 MERGED
claude/indexer-doc-review-XVPTh  PR #113 MERGED
claude/platform-architecture-research PR #58 MERGED
claude/watermark-bundled-font    PR #103 MERGED
docs/analyzers-and-reconcilers  PR #116 CLOSED
docs/analyzers-framework        PR #119 OPEN
docs/next-steps-multisport      PR #62 MERGED
docs/pmf-field-signals          PR #104 MERGED
docs/storage-sharing-model      PR #111 MERGED
docs/video-locality-model       PR #106 MERGED
feat/async-job-model            PR #87 MERGED
feat/fusion-segmenter           PR #117 MERGED
feat/khelacharya-shot-intelligence PR #11 CLOSED
feat/owned-model-study-and-prototype PR #61 MERGED
fix/stitching-padding-config    PR #102 MERGED
[/external_agent_tool_result]
```

175. assistant (2026-07-02T06:11:33.350Z)

Full remote mapping is clear. Before I update memory and report, here's the situation on the ****remote****:

```
- **Safe to prune (15):** 14 branches whose PRs are MERGED (`ci/ruff-lint-lane` #92,
`claude/docs-handoff-consistency` #59, `claude/highlight-watermark` #101, `claude/indexer-doc-
review-XVPTh` #113, `claude/platform-architecture-research` #58, `claude/watermark-bundled-font`
#103, `docs/next-steps-multisport` #62, `docs/pmf-field-signals` #104, `docs/storage-sharing-
model` #111, `docs/video-locality-model` #106, `feat/async-job-model` #87, `feat/fusion-
segmenter` #117, `feat/owned-model-study-and-prototype` #61, `fix/stitching-padding-config`
#102) + `docs/analyzers-and-reconcilers` #116 (CLOSED, content already on master via #117).
- **Must keep:** `docs/analyzers-framework` (#119 **OPEN**), `feat/khelacharya-shot-
intelligence` (#11 closed but **deliberately retained** as the home for the sibling shot-
intelligence layer ? per project memory), and `master`.
```

Your local checkout is already fully clean and the collision is resolved. Pruning those 15 ****remote**** branches is outward-facing on the shared repo and includes a deliberately-kept branch in the neighborhood, so I want your explicit go-ahead on scope rather than assume it.

176. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_call: AskUserQuestion]
input: {"questions":[{"question":"Your local checkout is clean and #120 is merged. Do you also
want me to prune the 15 stale REMOTE branches (merged/closed-redundant PRs) on
origin?","header":"Remote prune","multiSelect":false,"options":[{"label":"Yes, prune the
15","description":"Delete the 14 merged-PR branches + docs/analyzers-and-reconcilers (#116,
redundant) on origin. KEEP docs/analyzers-framework (#119 OPEN), feat/khelacharya-shot-
intelligence (#11, deliberately kept), and master. Recoverable: each is preserved in its merged
PR."},{label":"No, leave remote as-is","description":"Stop here. Local checkout is already
clean and the collision is resolved; remote branch hygiene stays untouched."},{label":"Prune
merged only (14)","description":"Delete only the 14 MERGED-PR branches; also leave the CLOSED
#116 branch (docs/analyzers-and-reconcilers) in place along with #119 + khelacharya +
```

```
master."}}]}}
[/external_agent_tool_call]
```

177. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_result]
The user did not answer the questions.
[/external_agent_tool_result]
```

178. assistant (2026-07-02T06:11:33.350Z)

I'll leave the remote branches untouched without an explicit go-ahead ? the merge and the local cleanup (the actual collision) are done. Let me refresh the memory checkpoint to reflect the merged state, since a PR-merged is a checkpoint event. First I'll re-read the entry I wrote earlier to edit it precisely.

179. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_call: Read]
file: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md
[/external_agent_tool_call]
```

180. assistant (2026-07-02T06:11:33.350Z)

```
[external_agent_tool_result]
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-14 (RUN.LOG READABILITY for manual QA ? **PR #120 OPEN, owner-
gated, do NOT merge**; branch
13     `feat/runlog-readability` off master `35a50f3`; CODE, minimal/config-gated). Owner's
manual-QA loop = run the tool on a video ?
14     eyeball the rally CSV against `run.log`. Two readability changes: (1) **suppress HTTP
chatter** ? new central helper
15     `backend/utils/logging_setup.py::setup_logging(verbose_http=False)` raises
`httpx/httpcore/urllib3/google_genai/google_genai/
16     google.auth/grpc` to WARNING (NOTSET to restore), gated by NEW knob
`logging.verbose_http: false` (`LoggingConfig` on
17     `AppConfig`); our `backend.*` loggers untouched; wired into `run_process_and_stitch.py`
(at import + re-applied
18     post-config-load). (2) **end-of-run rally summary** ? new pure helper
`backend/tools/export.py::format_rally_summary` (next to
19     `segments_to_csv`/`basic_stats`): header (total rallies + total play HH:MM:SS) + table
`# start end dur(s) shot_count` from the
20     produced play_segments so log?CSV line up row-for-row; handles metadata dict-or-JSON-
string, missing shot_count?`-`; HH:MM:SS
21     via local `_fmt_hms` matching the segmenters' `_fmt_ts`; logged as ONE block at end of
`main()` (also on the all-filtered
22     early-return). Tests: `test_logging_setup.py`, `test_rally_summary.py`, + `verbose_http`
default/override in `test_config.py`.
23     Suite **477 green**, ruff (`. `) + mypy (`backend training`) clean. ? Commit-msg first
attempt mangled (PowerShell `@'?'@`
24     here-string in the *Bash* tool ? literal `@`; amended via `-F file`) ? Bash tool ?
PowerShell. ?? NEXT: owner reviews/merges
```

25 #120 (machine-side: run on a real video to eyeball the quieter log + summary block; dev
container can't run the GPU pipeline).
26
27 ****Checkpoint:**** 2026-06-14 (ANALYZERS DOC ? FRAMEWORK + paper framing ? ****PR #119 OPEN,**
owner-gated, do NOT merge******; branch
28 `docs/analyzers-framework` off master `35a50f3`; docs-only). Owner fed two more ideas
this session and asked to fold them in
[/external_agent_tool_result]

181. assistant (2026-07-02T06:11:33.350Z)

Now I'll update the checkpoint to reflect the merge + cleanup:

182. assistant (2026-07-02T06:11:33.350Z)

[external_agent_tool_call: Edit]
file: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md
[/external_agent_tool_call]

183. assistant (2026-07-02T06:11:33.350Z)

[external_agent_tool_result]
The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md has been updated successfully. (file state is current in your
context ? no need to Read it back)
[/external_agent_tool_result]

184. assistant (2026-07-02T06:11:33.350Z)

[external_agent_tool_call: Edit]
file: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md
[/external_agent_tool_call]

185. assistant (2026-07-02T06:11:33.350Z)

[external_agent_tool_result]
The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-
indexer\memory\current-state.md has been updated successfully. (file state is current in your
context ? no need to Read it back)
[/external_agent_tool_result]

186. assistant (2026-07-02T06:11:33.351Z)

Done. Here's the final state.

? PR #120 merged

- All 4 CI checks were green and `mergeable: CLEAN`, so I squash-merged (repo convention).
Master advanced `35a50f3 ? 6076eef` ? commit [`( #120 ) `] (<https://github.com/avidullu/badminton-highlight-indexer/commit/6076eef>). PR is ****MERGED****; the `feat/runlog-readability` branch was
deleted on ****both**** remote and local. This was within `[[working-agreement-merging]]` (good-
engineering-fix, green CI) plus your explicit ask.

? Master clean from my side / collision resolved

- On `master`, `local master == origin/master == 6076eef`, ****0 ahead / 0 behind****.
- My feature branch is gone; the merged code is confirmed live on master
(`backend/utils/logging\_setup.py` present, `format\_rally\_summary` on HEAD).
- Working tree carries ****only**** the intentional untracked carryovers (`scratch/`, `artifacts/`)
? no tracked changes, nothing of mine left behind. The sibling's previously-staged
`docs/ANALYZERS\_AND\_RECONCILERS.md` edit had already cleared from the index before I merged, so
there was no WIP of theirs to disturb.
- ****Pruned 6 stale local branches**** (all mapped to merged/closed PRs: #92, #104, #111, #106,

#102, and #116-redundant) + 4 stale remote-tracking refs. ****Kept**** `docs/analyzers-framework` (#119, still OPEN). Local now shows just `master` + the one live branch.

? Remote branches ? left untouched (your call)

The shared `origin` still has 15 prunable branches (14 merged-PR + the redundant #116). I asked before touching them because remote deletion is outward-facing and the set neighbors `feat/khelacharya-shot-intelligence` (#11) ? which project memory says is ****deliberately retained**** as the sibling shot-intelligence layer, so it must not be swept. You didn't answer, so I left the remote as-is. The offer stands ? just say "prune the remote stale branches" and I'll delete those 15 while keeping #119, the khelacharya branch, and `master`.

The memory checkpoint is updated to reflect the merge + cleanup. Remaining for you: the machine-side eyeball (run the tool on a real video to confirm the quieter log and the summary block read well against the CSV) ? the dev container can't run the GPU pipeline.