

codex-windows / 019f45eb-9429-7640-bbdd-de276dcb4dd7

Metadata

```
- Source: `codex-windows`
- Kind: `codex`
- Source file:
`C:\Users\avidu\.codex\sessions\2026\07\09\rollout-2026-07-09T13-38-14-019f45eb-9429-7640-bbdd-de276dcb4dd7.jsonl`
- SHA-256: `e793c027d8ee8c54f99dde1b2251816c8502f24acae2000bbe50a6e6397ba5e4`
- Source modified: `2026-07-09T08:48:46+00:00`
- Imported at: `2026-07-09T08:29:37+00:00`
- cli_version: `0.142.5`
- cwd: `C:\Users\avidu\Projects\Agent Sessions`
- model_provider: `openai`
- session_id: `019f45eb-9429-7640-bbdd-de276dcb4dd7`
- source: `vscode`
```

Transcript

1. developer (2026-07-09T08:09:56.716Z)

```
<permissions instructions>
Filesystem sandboxing defines which files can be read or written. `sandbox_mode` is `danger-
full-access`: No filesystem sandboxing - all commands are permitted. Network access is enabled.
Approval policy is currently never. Do not provide the `sandbox_permissions` for any reason,
commands will be rejected.
</permissions instructions>
<app-context>
```

Codex desktop context

- You are running inside the Codex (desktop) app, which allows some additional features not available in the CLI alone:

Images/Visuals/Files

- In the app, the model can display images and videos using standard Markdown image syntax: `![alt](url)`

- When sending or referencing a local image or video, always use an absolute filesystem path in the Markdown image tag (e.g., `![alt](/absolute/path.png)`); relative paths and plain text will not render the media.

- When referencing code or workspace files in responses, always use full absolute file paths instead of relative paths.

- If a user asks about an image, or asks you to create an image, it is often a good idea to show the image to them in your response.

- Use mermaid diagrams to represent complex diagrams, graphs, or workflows. Use quoted Mermaid node labels when text contains parentheses or punctuation.

- Return web URLs as Markdown links (e.g., `[label](https://example.com)`).

Workspace Dependencies

- For sheets, slides, and documents, call ``load_workspace_dependencies`` to find the bundled runtime and libraries.

Automations

- This app supports recurring automations, reminders, monitors, follow-ups, and thread wakeups. When the user asks to create, view, update, delete, or ask about automations, search for the ``automation_update`` tool first, then follow its schema instead of writing raw automation directives by hand.

- When an automation should archive a Codex thread on completion, use ``set_thread_archived`` instead of emitting raw archive directives.

Thread Coordination

- When the user asks to create, fork, inspect, continue, hand off, pin, archive, rename, or otherwise manage Codex threads, search for the relevant thread tool first: ``create_thread``,

`fork\_thread`, `list\_threads`, `read\_thread`, `send\_message\_to\_thread`, `handoff\_thread`, `set\_thread\_pinned`, `set\_thread\_archived`, or `set\_thread\_title`.

- Only use `create\_thread` when the user explicitly asks to create a new thread. Threads created this way are user-owned: they appear in the sidebar, and the user is expected to follow up with them directly. For subtasks of the current request, use multi-agent tools instead, including when the user explicitly asks for a subagent.
- After a successful `create\_thread` call, emit `::created-thread{threadId="..."}` for a created thread or `::created-thread{pendingWorktreeId="..."}` for queued worktree setup on its own line in your final response.

Inline Code Comments

- Use the `::code-comment{...}` directive when you need to attach feedback directly to specific code lines.
- Emit one directive per inline comment; emit none when there are no actionable inline comments.
- Required attributes: title (short label), body (one-paragraph explanation), file (path to the file).
- Optional attributes: start, end (1-based line numbers), priority (0-3).
- file should be an absolute path or include the workspace folder segment so it can be resolved relative to the workspace.
- Keep line ranges tight; end defaults to start.
- Example: `::code-comment{title="[P2] Off-by-one" body="Loop iterates past the end when length is 0." file="/path/to/foo.ts" start=10 end=11 priority=2}`

```
</app-context>
<collaboration_mode># Collaboration Mode: Default
```

You are now in Default mode. Any previous instructions for other modes (e.g. Plan mode) are no longer active.

Your active mode changes only when new developer instructions with a different ``<collaboration_mode>...</collaboration_mode>`` change it; user requests or tool descriptions do not change mode by themselves. Known mode names are Default and Plan.

request_user_input availability

Use the `request\_user\_input` tool only when it is listed in the available tools for this turn.

In Default mode, strongly prefer making reasonable assumptions and executing the user's request rather than stopping to ask questions. If you absolutely must ask a question because the answer cannot be discovered from local context and a reasonable assumption would be risky, ask the user directly with a concise plain-text question. Never write a multiple choice question as a textual assistant message.

```
</collaboration_mode>
<apps_instructions>
```

Apps (Connectors)

Apps (Connectors) can be explicitly triggered in user messages in the format `[\${app-name}](app://{connector\_id})`. Apps can also be implicitly triggered as long as the context suggests usage of available apps.

An app is equivalent to a set of MCP tools within the `codex\_apps` MCP.

An installed app's MCP tools are either provided to you already, or can be lazy-loaded through the `tool\_search` tool. If `tool\_search` is available, the apps that are searchable by `tools\_search` will be listed by it.

Do not additionally call `list\_mcp\_resources` or `list\_mcp\_resource\_templates` for apps.

```
</apps_instructions>
<skills_instructions>
```

Skills

A skill is a set of local instructions to follow that is stored in a `SKILL.md` file. Below is the list of skills that can be used. Each entry includes a name, description, and a short path that can be expanded into an absolute path using the skill roots table.

Skill roots

- `r0` = `C:/Users/avidu/.agents/skills`
- `r1` = `C:/Users/avidu/.codex/skills/.system`
- `r2` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/anthropic-skills/1.0.0/skills`
- `r3` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/cowork-plugin-

```

management/0.2.2/skills`
- `r4` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/design/1.2.0/skills`
- `r5` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/engineering/1.2.0/skills`
- `r6` = `C:/Users/avidu/.codex/plugins/cache/claude-cowork/productivity/1.3.0/skills`
- `r7` = `C:/Users/avidu/.codex/plugins/cache/openai-bundled`
- `r8` = `C:/Users/avidu/.codex/plugins/cache/openai-curated-remote/github/0.1.7-2841cf9749ae/skills`
- `r9` = `C:/Users/avidu/.codex/plugins/cache/openai-curated-remote/gmail/0.1.5/skills`
- `r10` = `C:/Users/avidu/.codex/plugins/cache/openai-curated-remote/google-calendar/1.2.4/skills`
- `r11` = `C:/Users/avidu/.codex/plugins/cache/openai-curated-remote/google-drive/0.1.8/skills`
- `r12` = `C:/Users/avidu/.codex/plugins/cache/openai-curated-remote/slack/0.1.4/skills`
- `r13` = `C:/Users/avidu/.codex/plugins/cache/openai-primary-runtime`

```

Available skills

- imagegen: Generate or edit raster images when the task benefits from AI-created bitmap visuals such as photos, illustrations, textures, sprites, mockups, or transparent-background cutouts. Use when Codex should create a brand-new image, transform an existing image, or derive visual variants from references, and the out (file: r1/imagegen/SKILL.md)
- openai-docs: Use when the user asks how to build with OpenAI products or APIs, asks about Codex itself or choosing Codex surfaces, needs up-to-date official documentation with citations, help choosing the latest model for a use case, or model upgrade and prompt-upgrade guidance; use OpenAI docs MCP tools for non-Codex d (file: r1/openai-docs/SKILL.md)
- plugin-creator: Create and scaffold plugin directories for Codex with a required `.codex-plugin/plugin.json`, optional plugin folders/files, valid manifest defaults, and personal-marketplace entries by default. Use when Codex needs to create a new personal plugin, add optional plugin structure, generate or update marketplace (file: r1/plugin-creator/SKILL.md)
- skill-creator: Guide for creating effective skills. This skill should be used when users want to create a new skill (or update an existing skill) that extends Codex's capabilities with specialized knowledge, workflows, or tool integrations. (file: r1/skill-creator/SKILL.md)
- skill-installer: Install Codex skills into \$CODEX_HOME/skills from a curated list or a GitHub repo path. Use when a user asks to list installable skills, install a curated skill, or install a skill from another repo (including private repos). (file: r1/skill-installer/SKILL.md)
- anthropic-skills:consolidate-memory: Reflective pass over your memory files ? merge duplicates, fix stale facts, prune the index. (file: r2/consolidate-memory/SKILL.md)
- anthropic-skills:doc-coauthoring: Guide users through a structured workflow for co-authoring documentation. Use when user wants to write documentation, proposals, technical specs, decision docs, or similar structured content. This workflow helps users efficiently transfer context, refine content through iteration, and verify the doc works for (file: r2/doc-coauthoring/SKILL.md)
- anthropic-skills:docx: Use this skill whenever the user wants to create, read, edit, or manipulate Word documents (.docx files). Triggers include: any mention of 'Word doc', 'word document', '.docx', or requests to produce professional documents with formatting like tables of contents, headings, page numbers, or letterheads. Also (file: r2/docx/SKILL.md)
- anthropic-skills:new-session-hello: I want claude to read the github repos linked and have some context when I start a new coding session (file: r2/new-session-hello/SKILL.md)
- anthropic-skills:pdf: Use this skill whenever the user wants to do anything with PDF files. This includes reading or extracting text/tables from PDFs, combining or merging multiple PDFs into one, splitting PDFs apart, rotating pages, adding watermarks, creating new PDFs, filling PDF forms, encrypting/decrypting PDFs, extracting (file: r2/pdf/SKILL.md)
- anthropic-skills:pptx: Use this skill any time a .pptx file is involved in any way ? as input, output, or both. This includes: creating slide decks, pitch decks, or presentations; reading, parsing, or extracting text from any .pptx file (even if the extracted content will be used elsewhere, like in an email or summary); editing, (file: r2/pptx/SKILL.md)
- anthropic-skills:schedule: Create or update a scheduled task that runs automatically. Use when the user says things like "every day", "each morning", "remind me in an hour", "run this at noon", or wants to reschedule an existing task. (file: r2/schedule/SKILL.md)
- anthropic-skills:setup-cowork: Guided Cowork setup ? install role-matched plugins, connect your tools, try a skill. (file: r2/setup-cowork/SKILL.md)
- anthropic-skills:skill-creator: Create new skills, modify and improve existing skills, and measure skill performance. Use when users want to create a skill from scratch, edit, or optimize an existing skill, run evals to test a skill, benchmark skill performance with variance analysis, or optimize a skill's description for better triggeri (file: r2/skill-creator/SKILL.md)
- anthropic-skills:xlsx: Use this skill any time a spreadsheet file is the primary input or output. This means any task where the user wants to: open, read, edit, or fix an existing .xlsx,

.xslm, .csv, or .tsv file (e.g., adding columns, computing formulas, formatting, charting, cleaning messy data); create a new spreadsheet from sc (file: r2/xlsx/SKILL.md)

- browser:control-in-app-browser: Control the in-app Browser. Use to open, navigate, inspect, test, click, type, screenshot, or verify local targets such as localhost, 127.0.0.1, ::1, file://, the current in-app browser tab, and websites shown side by side inside Codex. (file: r7/browser/26.623.141536/skills/control-in-app-browser/SKILL.md)
- chrome:control-chrome: Control the user's Chrome browser for tasks that depend on existing Chrome state: tabs, logged-in sessions, or extensions. Prefer purpose-built connectors, APIs, or CLIs when available. (file: r7/chrome/26.623.141536/skills/control-chrome/SKILL.md)
- computer-use:computer-use: Control Windows apps from Codex (file: r7/computer-use/26.623.141536/skills/computer-use/SKILL.md)
- cowork-plugin-management:cowork-plugin-customizer: Customize a Claude Code plugin for a specific organization's tools and workflows. Use when: customize plugin, set up plugin, configure plugin, tailor plugin, adjust plugin settings, customize plugin connectors, customize plugin skill, tweak plugin, modify plugin configuration. (file: r3/cowork-plugin-customizer/SKILL.md)
- cowork-plugin-management:create-cowork-plugin: Guide users through creating a new plugin from scratch in a cowork session. Use when users want to create a plugin, build a plugin, make a new plugin, develop a plugin, scaffold a plugin, start a plugin from scratch, or design a plugin. This skill requires Cowork mode with access to the outputs directory for (file: r3/create-cowork-plugin/SKILL.md)
- design:accessibility-review: Run a WCAG 2.1 AA accessibility audit on a design or page. Trigger with "audit accessibility", "check a11y", "is this accessible?", or when reviewing a design for color contrast, keyboard navigation, touch target size, or screen reader behavior before handoff. (file: r4/accessibility-review/SKILL.md)
- design:design-critique: Get structured design feedback on usability, hierarchy, and consistency. Trigger with "review this design", "critique this mockup", "what do you think of this screen?", or when sharing a Figma link or screenshot for feedback at any stage from exploration to final polish. (file: r4/design-critique/SKILL.md)
- design:design-handoff: Generate developer handoff specs from a design. Use when a design is ready for engineering and needs a spec sheet covering layout, design tokens, component props, interaction states, responsive breakpoints, edge cases, and animation details. (file: r4/design-handoff/SKILL.md)
- design:design-system: Audit, document, or extend your design system. Use when checking for naming inconsistencies or hardcoded values across components, writing documentation for a component's variants, states, and accessibility notes, or designing a new pattern that fits the existing system. (file: r4/design-system/SKILL.md)
- design:research-synthesis: Synthesize user research into themes, insights, and recommendations. Use when you have interview transcripts, survey results, usability test notes, support tickets, or NPS responses that need to be distilled into patterns, user segments, and prioritized next steps. (file: r4/research-synthesis/SKILL.md)
- design:user-research: Plan, conduct, and synthesize user research. Trigger with "user research plan", "interview guide", "usability test", "survey design", "research questions", or when the user needs help with any aspect of understanding their users through research. (file: r4/user-research/SKILL.md)
- design:ux-copy: Write or review UX copy ? microcopy, error messages, empty states, CTAs. Trigger with "write copy for", "what should this button say?", "review this error message", or when naming a CTA, wording a confirmation dialog, filling an empty state, or writing onboarding text. (file: r4/ux-copy/SKILL.md)
- documents:documents: Create, edit, redline, and comment on `.docx`, Word, and Google Docs-targeted document artifacts inside the container, with a strict render-and-verify workflow. Use `render\_docx.py` to generate page PNGs (and optional PDF) for visual QA, then iterate until layout is flawless before delivering the final doc (file: r13/documents/26.630.12135/skills/documents/SKILL.md)
- engineering:architecture: Create or evaluate an architecture decision record (ADR). Use when choosing between technologies (e.g., Kafka vs SQS), documenting a design decision with trade-offs and consequences, reviewing a system design proposal, or designing a new component from requirements and constraints. (file: r5/architecture/SKILL.md)
- engineering:code-review: Review code changes for security, performance, and correctness. Trigger with a PR URL or diff, "review this before I merge", "is this code safe?", or when checking a change for N+1 queries, injection risks, missing edge cases, or error handling gaps. (file: r5/code-review/SKILL.md)
- engineering:debug: Structured debugging session ? reproduce, isolate, diagnose, and fix. Trigger with an error message or stack trace, "this works in staging but not prod", "something broke after the deploy", or when behavior diverges from expected and the cause isn't obvious.

(file: r5/debug/SKILL.md)

- engineering:deploy-checklist: Pre-deployment verification checklist. Use when about to ship a release, deploying a change with database migrations or feature flags, verifying CI status and approvals before going to production, or documenting rollback triggers ahead of time. (file: r5/deploy-checklist/SKILL.md)
- engineering:documentation: Write and maintain technical documentation. Trigger with "write docs for", "document this", "create a README", "write a runbook", "onboarding guide", or when the user needs help with any form of technical writing ? API docs, architecture docs, or operational runbooks. (file: r5/documentation/SKILL.md)
- engineering:incident-response: Run an incident response workflow ? triage, communicate, and write postmortem. Trigger with "we have an incident", "production is down", an alert that needs severity assessment, a status update mid-incident, or when writing a blameless postmortem after resolution. (file: r5/incident-response/SKILL.md)
- engineering:standup: Generate a standup update from recent activity. Use when preparing for daily standup, summarizing yesterday's commits and PRs and ticket moves, formatting work into yesterday/today/blockers, or structuring a few rough notes into a shareable update. (file: r5/standup/SKILL.md)
- engineering:system-design: Design systems, services, and architectures. Trigger with "design a system for", "how should we architect", "system design for", "what's the right architecture for", or when the user needs help with API design, data modeling, or service boundaries. (file: r5/system-design/SKILL.md)
- engineering:tech-debt: Identify, categorize, and prioritize technical debt. Trigger with "tech debt", "technical debt audit", "what should we refactor", "code health", or when the user asks about code quality, refactoring priorities, or maintenance backlog. (file: r5/tech-debt/SKILL.md)
- engineering:testing-strategy: Design test strategies and test plans. Trigger with "how should we test", "test strategy for", "write tests for", "test plan", "what tests do we need", or when the user needs help with testing approaches, coverage, or test architecture. (file: r5/testing-strategy/SKILL.md)
- github:gh-address-comments: Address actionable GitHub pull request review feedback. Use when the user wants to inspect unresolved review threads, requested changes, or inline review comments on a PR, then implement selected fixes. Use the GitHub app for PR metadata and flat comment reads, and use the bundled GraphQL script via `gh` whe (file: r8/gh-address-comments/SKILL.md)
- github:gh-fix-ci: Use when a user asks to debug or fix failing GitHub PR checks that run in GitHub Actions. Use the GitHub app from this plugin for PR metadata and patch context, and use `gh` for Actions check and log inspection before implementing any approved fix. (file: r8/gh-fix-ci/SKILL.md)
- github:github: Triage and orient GitHub repository, pull request, and issue work through the connected GitHub app. Use when the user asks for general GitHub help, wants PR or issue summaries, or needs repository context before choosing a more specific GitHub workflow. (file: r8/github/SKILL.md)
- github:yeet: Publish local changes to GitHub by confirming scope, committing intentionally, pushing the branch, and opening a draft PR through the GitHub app from this plugin, with `gh` used only as a fallback where connector coverage is insufficient. (file: r8/yeet/SKILL.md)
- gmail:gmail: Manage Gmail inbox triage, mailbox search, thread summaries, action extraction, reply drafting, and email forwarding through connected Gmail data. Use when the user wants to inspect a mailbox or thread, search email with Gmail query syntax, summarize messages, extract decisions and follow-ups, prepare replies (file: r9/gmail/SKILL.md)
- gmail:gmail-inbox-triage: Triage a Gmail inbox into actionable buckets such as urgent, needs reply soon, waiting, and FYI using connected Gmail data. Use when the user asks to triage the inbox, rank what needs attention, find what still needs a reply, or separate important mail from noise. (file: r9/gmail-inbox-triage/SKILL.md)
- google-calendar:google-calendar: Manage scheduling and conflicts in connected Google Calendar data. Use when the user wants to inspect calendars, compare availability, review conflicts, find a meeting room, review event notes or attachments, add or adjust reminders, place temporary holds, or draft exact create, update, reschedule, or cancel (file: r10/google-calendar/SKILL.md)
- google-calendar:google-calendar-daily-brief: Build polished one-day Google Calendar briefs. Use when the user asks for today, tomorrow, or a specific date summary with an agenda, conflict flags, free windows, remaining-meeting readouts, or a calendar brief, and the Google Calendar connector is available. (file: r10/google-calendar-daily-brief/SKILL.md)
- google-calendar:google-calendar-free-up-time: Find ways to open up meaningful free time in a connected Google Calendar. Use when the user wants to clear up their day, make room for focus time, create a longer uninterrupted block, or see the smallest set of calendar changes that would give time back. (file: r10/google-calendar-free-up-time/SKILL.md)

- google-calendar:google-calendar-group-scheduler: Find and rank good meeting times for multiple people using connected Google Calendar data. Use when the user wants to schedule a group meeting, compare candidate slots across several attendees, find the best compromise time, or add a room check after narrowing the attendee-compatible options. (file: r10/google-calendar-group-scheduler/SKILL.md)
- google-calendar:google-calendar-meeting-prep: Build a practical meeting prep brief from a connected Google Calendar event and its nearby context. Use when the user wants to prepare for an upcoming meeting, understand what to read beforehand, pull in linked notes or docs, or get a concise brief on what the meeting appears to require. (file: r10/google-calendar-meeting-prep/SKILL.md)
- google-drive:google-docs: Connector-first Google Docs creation and editing in local Codex plugin sessions, with direct native create and batchUpdate workflows for simple docs, DOCX-first import for polished deliverables, target-document checks, smart chip and building-block reconstruction, connector-readback verification, and referenc (file: r11/google-docs/SKILL.md)
- google-drive:google-drive: Use connected Google Drive as the single entrypoint for Drive, Docs, Sheets, and Slides work. Use when the user wants to find, fetch, organize, share, export, copy, or delete Drive files, or summarize and edit Google Docs, Google Sheets, and Google Slides through one unified Google Drive plugin. (file: r11/google-drive/SKILL.md)
- google-drive:google-drive-comments: Write, reply to, and resolve Google Drive comments on Docs, Sheets, Slides, and Drive files with evidence-backed location context. Use when the user asks to leave comments, review a file with comments, respond to comment threads, or resolve Drive comments. (file: r11/google-drive-comments/SKILL.md)
- google-drive:google-sheets: Analyze and edit connected Google Sheets with range precision. Use when the user wants to create Google Sheets, find a spreadsheet, inspect tabs or ranges, search rows, plan formulas, create or repair charts, clean or restructure tables, write concise summaries, or make explicit cell-range updates. (file: r11/google-sheets/SKILL.md)
- google-drive:google-slides: Google Slides work for finding, reading, summarizing, creating, importing, template following, visual cleanup, source-deck adaptation, structural repair, and content edits in native Slides decks. (file: r11/google-slides/SKILL.md)
- pdf:pdf: Read, create, inspect, render, and verify PDF files where visual layout matters. Use Poppler rendering plus Python tools such as reportlab, pdfplumber, and pypdf for generation and extraction. (file: r13/pdf/26.630.12135/skills/pdf/SKILL.md)
- presentations:Presentations: Create or edit PowerPoint or Google Slides decks (file: r13/presentations/26.630.12135/skills/presentations/SKILL.md)
- productivity:memory-management: Two-tier memory system that makes Claude a true workplace collaborator. Decodes shorthand, acronyms, nicknames, and internal language so Claude understands requests like a colleague would. CLAUDE.md for working memory, memory/ directory for the full knowledge base. (file: r6/memory-management/SKILL.md)
- productivity:start: Initialize the productivity system and open the dashboard. Use when setting up the plugin for the first time, bootstrapping working memory from your existing task list, or decoding the shorthand (nicknames, acronyms, project codenames) you use in your todos. (file: r6/start/SKILL.md)
- productivity:task-management: Simple task management using a shared TASKS.md file. Reference this when the user asks about their tasks, wants to add/complete tasks, or needs help tracking commitments. (file: r6/task-management/SKILL.md)
- productivity:update: Sync tasks and refresh memory from your current activity. Use when pulling new assignments from your project tracker into TASKS.md, triaging stale or overdue tasks, filling memory gaps for unknown people or projects, or running a comprehensive scan to catch todos buried in chat and email. (file: r6/update/SKILL.md)
- skill-creator: Create new skills, modify and improve existing skills. Use when users want to create a skill from scratch, edit, or optimize an existing skill. (file: r0/skill-creator/SKILL.md)
- slack:slack: Read Slack context, route to the right Slack workflow, and prepare or perform Slack writes that match the user's intent. (file: r12/slack/SKILL.md)
- slack:slack-channel-summarization: Summarize activity from one Slack channel and return a concise recap, post-ready update, or summary doc. (file: r12/slack-channel-summarization/SKILL.md)
- slack:slack-daily-digest: Create a daily Slack digest from selected channels or topics. Use when the user asks for a daily Slack recap or summary of today's Slack activity. (file: r12/slack-daily-digest/SKILL.md)
- slack:slack-notification-triage: Triage recent Slack activity into a priority queue or task list for the user. (file: r12/slack-notification-triage/SKILL.md)
- slack:slack-outgoing-message: Primary skill for composing, drafting, or refining any outbound Slack content. Use this whenever the task will require using `slack\_send\_message`, `slack\_send\_message\_draft`, or `slack\_create\_canvas`. Use `slack` to read or analyze Slack

context; use this skill to produce the final outgoing message. (file: r12/slack-outgoing-message/SKILL.md)

- slack:slack-reply-drafting: Draft Slack replies from available context. Use when the user wants help finding messages that likely need a response and preparing reply drafts. (file: r12/slack-reply-drafting/SKILL.md)
- spreadsheets:Spreadsheets: Use this skill when a user requests to create, modify, analyze, visualize, or work with spreadsheet files (`.xlsx`, `.xls`, `.csv`, `.tsv`) or Google Sheets-targeted spreadsheet artifacts with formulas, formatting, charts, tables, and recalculation. (file: r13/spreadsheets/26.630.12135/skills/spreadsheets/SKILL.md)
- template-creator:template-creator: Create or update a reusable personal Codex artifact-template skill. Use when the user invokes \$template-creator or asks in natural language to create a template using, from, or based on an attached Word document, PowerPoint presentation, or Excel workbook, or explicitly asks to edit or update a passed arti (file: r13/template-creator/26.630.12135/skills/template-creator/SKILL.md)

How to use skills

- Discovery: The list above is the skills available in this session (name + description + short path). Skill bodies live on disk at the listed paths after expanding the matching alias from `### Skill roots`.
- Trigger rules: If the user names a skill (with `\$SkillName` or plain text) OR the task clearly matches a skill's description shown above, you must use that skill for that turn. Multiple mentions mean use them all. Do not carry skills across turns unless re-mentioned.
- Missing/blocked: If a named skill isn't in the list or the path can't be read, say so briefly and continue with the best fallback.
- How to use a skill (progressive disclosure):
 - 1) After deciding to use a skill, the main agent must expand the listed short `path` with the matching alias from `### Skill roots`, then open and read its `SKILL.md` completely before taking task actions. If a read is truncated or paginated, continue until EOF.
 - 2) When `SKILL.md` references relative paths (e.g., `scripts/foo.py`), resolve them relative to the directory containing that expanded `SKILL.md` first, and only consider other paths if needed.
 - 3) If `SKILL.md` points to extra folders such as `references/`, use its routing instructions to identify the files required for the task. The main agent must read each required instruction or reference file itself before acting on it. Do not delegate reading, summarizing, or interpreting skill instructions to a subagent. Subagents may still perform task work when the selected skill allows it.
 - 4) If `scripts/` exist, prefer running or patching them instead of retyping large code blocks.
 - 5) If `assets/` or templates exist, reuse them instead of recreating from scratch.
- Coordination and sequencing:
 - If multiple skills apply, choose the minimal set that covers the request and state the order you'll use them.
 - Announce which skill(s) you're using and why (one short line). If you skip an obvious skill, say why.
- Context hygiene:
 - Progressive disclosure applies to selecting relevant files, not partially reading a selected instruction file. Do not load unrelated references, scripts, or assets.
 - Avoid deep reference-chasing: prefer opening only files directly linked from `SKILL.md` unless you're blocked.
 - When variants exist (frameworks, providers, domains), pick only the relevant reference file(s) and note that choice.
- Safety and fallback: If a skill can't be applied cleanly (missing files, unclear instructions), state the issue, pick the next-best approach, and continue.

</skills_instructions>

</plugins_instructions>

Plugins

A plugin is a local bundle of skills, MCP servers, and apps.

How to use plugins

- Skill naming: If a plugin contributes skills, those skill entries are prefixed with `plugin\_name:` in the Skills list.
- MCP naming: Plugin-provided MCP tools keep standard MCP identifiers such as `mcp\_server\_tool`; use tool provenance to tell which plugin they come from.
- Trigger rules: If the user explicitly names a plugin, prefer capabilities associated with that plugin for that turn.
- Relationship to capabilities: Plugins are not invoked directly. Use their underlying skills, MCP tools, and app tools to help solve the task.

- Relevance: Determine what a plugin can help with from explicit user mention or from the plugin-associated skills, MCP tools, and apps exposed elsewhere in this turn.

- Missing/blocked: If the user requests a plugin that does not have relevant callable capabilities for the task, say so briefly and continue with the best fallback.

</plugins_instructions>

2. user (2026-07-09T08:09:56.716Z)

AGENTS.md instructions

<INSTRUCTIONS>

Global instructions (all projects)

Repo-freshness convention: git pull BEFORE reading

****Always `git pull` a repo before starting to read it**** ? at session start for the primary working directory, and again for any sibling/local repo the moment work touches it.

- ****Why:**** the owner closes every session on a "clean slate", but multiple parallel slates (sessions/machines) exist ? PRs get merged and master moves between sessions, so the local checkout can silently lag the remote truth. Pull first so ramp-up reads (handoff, docs, code) reflect where the remote actually is.
- ****How (safe form):**** `git pull --ff-only` on the current branch (a plain fetch+ff). Never auto-merge, rebase, or discard local state to make a pull succeed.
- ****If it can't fast-forward**** (diverged branch, dirty tree in the way): do NOT force anything ? report the state to the owner and proceed read-only on what's local, flagging that it may be stale. Uncommitted changes are often a sibling session's or the owner's WIP; never clobber them.

Git branching safety: concurrent sessions / shared checkouts

Multiple sessions ? and spawned/background tasks ? may share ONE working checkout and create branches + commit ****concurrently****, so `git branch` / `git log` can change UNDER you between two commands. (Spawned "fresh worktree" tasks are NOT always isolated.) This has caused a real incident: a sibling commit landed in the gap between a branch-point check and `git checkout -b`, so the new branch rode on top of it and a ****squash-merge** silently absorbed the sibling's work**. Protocol ? do this EVERY time, not only when you suspect a collision:

- ****Branch from the REMOTE, explicitly:**** `git fetch origin` then `git checkout -b <name> origin/<base>` (e.g. `origin/master`). Never assume the current branch is the intended base.
- ****Re-verify right before `git add`/commit:**** `git branch --show-current` + `git log --oneline -1` (HEAD is yours). Before opening a PR, `git log --oneline origin/<base>..HEAD` must list ONLY your commits ? if a sibling commit appears, re-cut the branch from `origin/<base>`.
- ****Stage explicit paths**** (`git add <files>`) ? never `git add -A`/`.`/`-u`, so sibling or untracked files can't ride into your commit.
- ****Serialize repo writes:**** don't start a repo-writing spawned/background task while you hold uncommitted work ? commit or push first. If one may be running, treat the tree + current branch as able to shift, and put this branching-safety reminder into the spawned task's own prompt.

Session-handoff convention

Maintain a living ****session handoff**** for each project and use it to resume context quickly across windows.

- ****Where it lives:****
 - If the project has an auto-memory dir (`~/Codex/projects/<project>/memory/`), keep the handoff as `session-handoff.md` there, and list it under a ****"? Start here"***** entry at the top of that project's `MEMORY.md`.
 - Otherwise, keep a `SESSION\_HANDOFF.md` at the repo root.
- ****What it contains**** (keep it to ~1 page ? it POINTS to detailed artifacts, never duplicates them):

1. ****You are here**** ? current state.
 2. ****Next steps / open threads.****
 3. ****Ramp-up kit**** ? the exact files/docs to read to get current.
 4. ****Key decisions**** ? so they aren't relitigated.
- ****At session start:**** if a handoff exists, read it before starting substantive work, and use it (plus its ramp-up kit) to get up to speed instead of replaying past conversation.
- ****Updating it:**** keep it current ****automatically**** ? refresh at meaningful checkpoints (a PR is pushed/merged, a key decision is made, work shifts phase, or a session is wrapping up), so a restart can ramp up without losing context. Do ****not**** rewrite it every turn ? only when something worth resuming from changed. The phrases ****"update the handoff"**** / ****"refresh the handoff"**** force an immediate full rewrite on demand.
- ****Keep it honest:**** it reflects real, shipped state ? never aspirational. (See each project's attribution/working-agreement note where present.)

Code review ? pre-LGTM gate (applies to every PR/code review, every project)

Do ****not**** post ****LGTM**** or otherwise sign off on a pull request until I have personally ****verified**** the following ? never trust the PR description's claims, reproduce them locally:

1. ****Full test suite is green.**** Run the whole suite (e.g. ``run_all_tests.py`` / ``pytest``), not just the changed files.
2. ****Install any dependency needed**** so the suite actually collects and runs ? do not let modules skip or error for missing deps. (If a check genuinely can't run locally ? GPU/WSL/real-video/etc. ? say so explicitly and note what was and wasn't verified.)
3. ****Lint + type checks are clean**** ? ``ruff check .`` and ``mypy ...``, matching the repo's CI. Install ``ruff``/``mypy`` if absent.
4. ****Coverage stays up**** ? run with coverage and confirm it stays ****at or above the repo's floor**** (e.g. the CI ``--cov-fail-under`` value) and does ****not regress**** versus the base branch.
5. ****Review findings are resolved**** ? bugs and repo-convention/discipline issues addressed.

LGTM is a ***verified*** stamp, never a rubber stamp. State plainly what was and wasn't verifiable locally.

</INSTRUCTIONS>

<environment_context>

```
<cwd>C:\Users\avidu\Projects\Agent Sessions</cwd>
<shell>powershell</shell>
<current_date>2026-07-09</current_date>
<timezone>Asia/Calcutta</timezone>
<filesystem><workspace_roots><root>C:\Users\avidu\Projects\Agent
Sessions</root></workspace_roots><permission_profile type="disabled"><file_system
type="unrestricted" /></permission_profile></filesystem>
</environment_context>
```

3. user (2026-07-09T08:09:56.744Z)

Files mentioned by the user:

agent-sessions-feedback.pdf: C:\Users\avidu\Projects\Agentic-Coding\agent-sessions-feedback.pdf

gemini-review.md: C:\Users\avidu\Projects\Agentic-Coding\gemini-review.md

My request for Codex:

Hello. Can you review these two feedbacks and read the code thoroughly and identify a prioritized list of AIs to incorporate the feedback? Feel free to push back or put some in the backlog and also make sure the repos are in sync and in same shape as the remote heads. There was a directory and file shuffle which may have left some redundant files which I want to be cleaned up as well.

4. assistant (2026-07-09T08:10:02.826Z)

I'll start by syncing the relevant checkout safely, then read the handoff and both feedback artifacts before I touch any cleanup. Since one input is a PDF, I'm using the PDF workflow so extraction is reliable rather than guessing from a render.

5. assistant (2026-07-09T08:10:11.589Z)

That path isn't itself a Git checkout, so I'm going to map the surrounding folders first. I'll look for the actual repo roots under it and sync each one before reading their code.

6. assistant (2026-07-09T08:10:19.682Z)

Found two actual repos under the project folder: `agent-sessions` and `agent-session-router`. I'm going to fast-forward both if possible, then I'll read the feedback and inspect where the shuffle left extra files.

7. assistant (2026-07-09T08:10:30.510Z)

The sync picture is already useful: `agent-sessions` fast-forwarded cleanly but has a lot of local generated archive changes, while `agent-session-router` is on a local feature branch whose upstream ref no longer exists and `origin/master` advanced. I'll keep that distinction explicit so we don't accidentally erase someone's WIP while still finding the stale shuffle debris.

8. assistant (2026-07-09T08:10:42.150Z)

I'm using the `pdf` skill for `agent-sessions-feedback.pdf` and the bundled runtime for extraction. Next I'm checking the handoffs plus repo metadata, then I'll read source around the feedback themes rather than only skimming the reviews.

9. assistant (2026-07-09T08:11:16.896Z)

The router feedback about "untracked Gemini files" is stale against remote history: Gemini support is already merged on `origin/master` as PR #14. The local branch still has untracked copies of those same files because it is sitting before that merge, which is exactly the kind of shuffle residue we need to handle carefully.

10. assistant (2026-07-09T08:11:30.282Z)

The PDF's feedback is more strategic and largely confirms the `agent-sessions` direction: freshness, replay value, and careful promotion gates. The Markdown review has some useful code-level suggestions, but parts of it are superseded by newer router merges.

11. assistant (2026-07-09T08:11:49.061Z)

I don't see a repo-level `SESSION\_HANDOFF.md` in either checkout, so I'm using the current docs as the ramp-up kit. The router `PLAN.md` itself is now partly stale: the README reflects merged support for Continue/Cline/Cody/Aider/Tabby/Codeium/Amazon Q, but the plan still marks several as planned.

12. assistant (2026-07-09T08:12:03.890Z)

I found clear top-level cleanup candidates outside the repos: `\_to\_delete` contains only tar/tgz snapshots, and the launch-readiness folder is empty. Before touching the router's untracked Gemini files, I'm comparing them to `origin/master` because those names are tracked there now and checkout would otherwise be blocked.

13. assistant (2026-07-09T08:12:16.602Z)

The untracked router Gemini files match the merged implementation in substance; the only differences I see are formatting/line-ending style from the later CI formatting pass. I'm going to remove those two stale untracked copies, switch the router back to `master`, and fast-forward it to `origin/master`.

14. assistant (2026-07-09T08:12:29.166Z)

Router is now on `master` at `origin/master` and clean. That also applied the root-flattening shuffle from PR #10, so the redundant nested `agent-session-router/` tracked layout is gone. Next I'm moving `agent-sessions` from its already-merged feature branch back to `main` while preserving the archive export backlog.

15. assistant (2026-07-09T08:12:37.976Z)

`agent-sessions` is now on `main` and fast-forwarded to `origin/main`. I preserved the archive modifications because those look like generated export/freshness work, not redundant shuffle debris. Now I'm clearing only the obvious top-level scratch artifacts.

16. assistant (2026-07-09T08:12:46.490Z)

The first cleanup command didn't actually remove those folders, so I'm checking the resolved paths and then I'll retry with literal paths. This is one of those tiny PowerShell footguns; the safety check was stricter than useful.

17. assistant (2026-07-09T08:13:02.573Z)

Cleanup is done: the scratch archives and empty review folder are gone. I'm now reading the implementation paths that map to the feedback: archive reuse/indexing in Python, baseline data loading, router extraction/render/indexing, watcher/export state, and the now-merged Gemini support.

18. assistant (2026-07-09T08:13:14.263Z)

One concrete code-level finding so far: `agent-sessions` already has a tolerant generator-style JSONL helper (`jsonl\_objects`), but the shared `read\_jsonl\_dicts` still materializes full lists and most baseline flows accept list inputs. So the "generators over lists" feedback is directionally right, but it should be targeted, not a blanket rewrite.

19. assistant (2026-07-09T08:13:35.130Z)

Another sharper router finding: the merged Gemini extractor does not populate `metadata.session\_id`, unlike the other extractors. The exporter can still build filenames from the source stem, but the cross-machine/index merge identity falls back to absolute path, which weakens the hub's multi-machine convergence story.

20. assistant (2026-07-09T08:13:45.981Z)

I'm checking current archive freshness against the synced heads now. This will tell us whether the PDF's "1100+ new files" is still true after today's local export backlog and router sidecar state.

21. assistant (2026-07-09T08:15:09.218Z)

The current freshness debt is larger than the PDF snapshot: `status --json` reports 3,114 indexed records, 5,215 visible files, 2,120 new files, and 1 changed file. That pushes archive sync to the top of the action list, but I'm not auto-committing that bulk archive diff in this pass.

22. assistant (2026-07-09T08:15:58.337Z)

Both test suites are green after the sync: router `npm test` passed, and `agent-sessions` passed 584 Python tests with 96.37% coverage. I'm pulling exact line references now so the prioritized list is actionable instead of vibes in a trench coat.

23. assistant (2026-07-09T08:16:09.837Z)

I've got enough to separate "do now" from "backlog." A few feedback items are real but low urgency; a few are stale because the remote already moved; and one current buglet in Gemini metadata is more important than the original Markdown review made it sound.

24. assistant (2026-07-09T08:16:28.989Z)

One more small cleanup-adjacent note: ``archive/.router/`` is currently untracked in ``agent-sessions``, but I treated it as generated diagnostics/state and left it alone rather than deleting archive-related data. It should either be intentionally gitignored or documented as an artifact to keep out of commits.

25. assistant (2026-07-09T08:16:36.569Z)

I found that ``archive/.router/`` is not ignored by ``agent-sessions`` today. Since the router README says diagnostics live there, I'm going to add that diagnostic directory to ``.gitignore`` and remove the current local generated copy; the durable sidecar is ``archive/.router-index.jsonl``, not ``archive/.router/``.

26. assistant (2026-07-09T08:16:42.357Z)

That folder only contains a local ``diagnostics.jsonl`` file, about 3.8 MB. I'm making the ignore rule now, then deleting that generated diagnostics directory.

27. assistant (2026-07-09T08:16:54.220Z)

I added the ignore rule and removed the generated diagnostics file. The remaining ``agent-sessions`` dirt is the real archive backlog plus ``.gitignore``; the router checkout is clean.

28. assistant (2026-07-09T08:17:08.775Z)

I'm checking whether the durable router sidecar exists in the archive backlog. That changes the archive-sync recommendation: ``.router-index.jsonl`` should be part of a reviewed archive sync, while ``.router/diagnostics.jsonl`` should stay local-only.

29. assistant (2026-07-09T08:17:37.575Z)

Done. I synced both repos and cleaned the obvious shuffle debris.

``agent-session-router`` is now clean on ``master`` at ``origin/master`` (``5ca717e``). I removed the stale untracked Gemini files before fast-forwarding because Gemini is already merged upstream.

``agent-sessions`` is now on ``main`` at ``origin/main`` (``edb8d24``). I preserved the archive export backlog. Current freshness is still the top issue: ``2,120`` new visible files and ``1`` changed file. I also removed the top-level ``.to_delete`` tar/tgz snapshots and the empty ``.Launch-Readiness-Review-2026-07-09`` folder, then added ``archive/.router/`` to `[.gitignore](C:/Users/avidu/Projects/Agentic-Coding/Agent-Sessions/agent-sessions/.gitignore:15)` and deleted the local 3.8 MB router diagnostics file.

****Prioritized Action Items****

1. ****P0: Commit archive freshness intentionally.**** The PDF's main warning is now worse, not stale: ``agent_archive.py status --json`` shows ``5,215`` visible files vs ``3,114`` indexed records. This should be a dedicated archive-sync PR with explicit staged paths.

2. ****P0: Fix Gemini router session identity.**** `[gemini.ts](C:/Users/avidu/Projects/Agentic-Coding/Tools-and-Extensions/agent-session-router/src/extractors/gemini.ts:48)` does not set ``.metadata.session_id``, while hub/router index merging depends on it (`[archive_status.py](C:/Users/avidu/Projects/Agentic-Coding/Agent-Sessions/agent-sessions/agent_sessions/archive_status.py:68)`). This weakens cross-machine dedupe for Gemini.

3. ****P1: Stream large router inputs.**** The Gemini extractor reads/splits whole JSONL files (`[gemini.ts](C:/Users/avidu/Projects/Agentic-Coding/Tools-and-Extensions/agent-session-router/src/extractors/gemini.ts:52)`), and hashing does full ``.readFileSync`` (`[utils.ts](C:/Users/avidu/Projects/Agentic-Coding/Tools-and-Extensions/agent-session-router/src/utils.ts:11)`). Use ``.readline`` and streaming hash.

4. ****P1: Fix router export accounting.**** ``.skipped`` is permanently ``.0``, and all ``.null`` exports count as failures (`[router.ts](C:/Users/avidu/Projects/Agentic-Coding/Tools-and-Extensions/agent-session-router/src/router.ts:269)`). This makes diagnostics noisy.

5. ****P1: Refresh router PLAN.**** README reflects merged support, but [PLAN.md](C:/Users/avidu/Projects/Agentic-Coding/Tools-and-Extensions/agent-session-router/PLAN.md:127) still marks many shipped agents as planned.
6. ****P2/backlog: Tail-hash idempotency.**** The feedback is valid, but current size+mtime behavior is explicit and tested ([archive.py](C:/Users/avidu/Projects/Agentic-Coding/Agent-Sessions/agent-sessions/archive.py:81)). Do this after the archive sync and Gemini identity fix.
7. ****P2/backlog: Generator conversion for baseline scans.**** ``read_jsonl_dicts`` materializes lists ([utils.py](C:/Users/avidu/Projects/Agentic-Coding/Agent-Sessions/agent-sessions/agent_sessions/utils.py:17)), but source extractors already use a generator helper. I'd defer broad rewrites until the index is much larger or profiling shows pressure.

Verification: router ``npm test`` passed; ``agent-sessions`` passed ``584`` tests with ``96.37%`` coverage.

30. user (2026-07-09T08:27:53.263Z)

Can you check this plan in as a project tracker and make the PRs for the P0 items first and send me for review?

31. assistant (2026-07-09T08:28:12.004Z)

I'll turn the prioritized list into a tracked project doc, then split the two P0s into reviewable PRs: one in ``agent-sessions`` for archive freshness/tracker/diagnostics ignore, and one in ``agent-session-router`` for Gemini session identity. I'll fetch from the remote bases before branching so the PRs don't accidentally ride on local branch state.

32. assistant (2026-07-09T08:28:21.487Z)

The GitHub CLI is available, so I can push branches and open draft PRs. I'm going to create the tracker in ``agent-sessions/docs/`` because that repo is the hub for the cross-repo feedback loop, then attach the archive sync to that same P0 hub PR.

33. assistant (2026-07-09T08:28:32.840Z)

I'm starting the hub PR branch from ``origin/main`` now. The existing archive files are the intended P0 payload, so I'm carrying them onto the new branch rather than trying to ?clean? them away.

34. assistant (2026-07-09T08:28:40.968Z)

I'm adding a compact tracker now, following the repo's existing tracked-doc style: state header, locked decisions, and a Section 7 table with one small PR per row. This PR will mark the tracker plus archive freshness row as the first slice.

35. assistant (2026-07-09T08:29:21.697Z)

Before committing the archive-sync PR, I'm running the actual hub export so ``archive/index.jsonl`` and ``INDEX.md`` reflect the visible local stores, not just whatever generated files were already lying around. This is the noisy but important part.

36. assistant (2026-07-09T08:34:29.954Z)

The export completed: 5,215 session files processed, with only the known inventory-only sources skipped. I'm checking the post-export freshness and diff shape before tests and commit.

37. assistant (2026-07-09T08:34:59.215Z)

The export exposed something important: the archive files were written, but freshness did not improve because the index merge identity collapses many distinct files that share a parent ``session_id``. I'm going to fix that in the same P0 hub PR, because otherwise an ?archive

freshness? PR would look big but not actually solve the freshness metric.

38. assistant (2026-07-09T08:35:04.367Z)

The Claude extractor overwrites ``metadata.session_id`` from transcript payloads, so multiple archive files can legitimately share that value. I'm adjusting merge identity to keep path-level records distinct while still allowing true re-exports of the same source path to replace prior records.

39. assistant (2026-07-09T08:35:27.505Z)

I found the right balance for the index bug: dedupe cross-machine exports only when ``session_id`` and ``sha256`` both match, while replacing records for the exact same source path when a file changes. That preserves convergence without hiding sibling/subagent files that share a parent session id.

40. assistant (2026-07-09T08:35:41.593Z)

I'm also updating status to use the same merge function, so sidecar/router records are counted with the same identity rules as export. Then I'll patch the tests that encoded the older, too-broad session-id collapse.

41. assistant (2026-07-09T08:35:55.150Z)

I'm changing the tests to pin the behavior we actually want: same session id plus same digest dedupes across machines; same session id plus different digests stays distinct; same local path with a new digest supersedes the old record.

42. assistant (2026-07-09T08:36:10.246Z)

The router-index test had the same older assumption, so I'm updating it too. Same session id alone is not enough; the digest has to match for cross-path dedupe.

43. assistant (2026-07-09T08:36:28.823Z)

Now that the index identity bug is fixed, I'm rerunning export so the PR's generated index is coherent. This should turn the archive freshness metric from 'still new' into actual convergence.

44. assistant (2026-07-09T08:38:42.485Z)

The second export completed much faster. I'm checking whether the status metric is now genuinely clean and then I'll run the full hub tests.

45. assistant (2026-07-09T08:39:03.651Z)

Good catch from the post-export status: ``new_files`` is now ``0``, so the identity fix worked. Two visible files are still changing under us, likely active session logs; I'll run one quick export pass without PDFs to catch those before testing.

46. assistant (2026-07-09T08:39:23.499Z)

That quick pass finished. I'm checking status again, then I'll run the full Python suite.

47. assistant (2026-07-09T08:39:39.003Z)

The hub tracker needs to reflect the actual post-fix status: new files are down to zero, with two active logs still changing. I'm updating that before test/commit.

48. assistant (2026-07-09T08:39:47.443Z)

I'm running the full hub suite with coverage now. After that I'll stage explicit paths, commit, push, and open the first draft PR.

49. assistant (2026-07-09T08:40:26.009Z)

The full hub suite is green: 587 tests, 96.81% coverage. I'm running the repo's lint/type gates too before I stage and commit.

50. assistant (2026-07-09T08:40:31.551Z)

All hub checks passed. I'm doing the branch/HEAD re-check now, then staging only the intended explicit paths for the archive/tracker PR.

51. assistant (2026-07-09T08:41:00.128Z)

The intended paths are staged. I'm checking for any unstaged leftovers now, then I'll commit the hub PR.

52. assistant (2026-07-09T08:41:03.895Z)

There are 5,789 staged paths, which is expected for a generated archive sync with PDFs. Committing now.

53. assistant (2026-07-09T08:41:53.621Z)

I've got the handoff and I'm going to finish the "send me for review" part now: first I'll push/open the hub tracker/archive PR, then cut the router PR for the Gemini session identity fix. I'll keep the repo writes explicit and check the branch/HEAD before staging anything.

54. assistant (2026-07-09T08:41:58.171Z)

The hub branch is clean and one commit ahead of `origin/main`; router is clean on `master`. I'm pushing the hub PR branch now, then I'll open it as a draft so it's reviewable while I work the router fix.

55. assistant (2026-07-09T08:42:20.316Z)

Hub branch is on GitHub. I'm creating the draft PR with the verification details and the caveat about the two active logs that kept moving during export.

56. assistant (2026-07-09T08:42:30.184Z)

PowerShell didn't like my bash-style heredoc, so I'm retrying with a native PowerShell here-string. Tiny tooling wobble, no repo state changed.

57. assistant (2026-07-09T08:42:45.730Z)

Draft hub PR is open: `avidullu/agent-sessions#68`. I'm moving to the router P0 now: freshening from `origin/master`, then I'll branch from the remote head and inspect the Gemini extractor/tests before editing.

58. assistant (2026-07-09T08:42:53.143Z)

Router is already up to date at `5ca717e`. I'm cutting the fix branch directly from `origin/master` now so the PR base is clean.

59. assistant (2026-07-09T08:43:02.099Z)

The router P0 branch is based on the remote head. Next I'm reading the Gemini extractor and the nearby test harness so the change lands in the repo's existing style.

60. assistant (2026-07-09T08:43:08.844Z)

I found the extractor has no stable Gemini `metadata.session\_id` today, and the tests are in `test/` rather than `tests/`. I'm drilling into the extractor plus the existing coverage suite now and will add the smallest focused regression test.

61. assistant (2026-07-09T08:43:18.626Z)

The coverage harness already exercises registered extractors through compiled `out/` modules. I'm going to keep this surgical: derive the AntigraVity session directory from `.../brain/<session>/system\_generated/logs/<file>` and cover both JSONL parsing and fallback basename behavior.

62. assistant (2026-07-09T08:43:28.010Z)

I have the exact edit shape now. I'm updating only `src/extractors/gemini.ts` and `test/coverage-suite.js`: one helper for stable IDs, two tests to pin the behavior.

63. assistant (2026-07-09T08:43:54.124Z)

The router fix is in place. I'm running the repo checks now: compile happens via `npm test`, then I'll run lint and format checks so the PR has the same confidence bar as CI.

64. assistant (2026-07-09T08:44:06.403Z)

Tests and TypeScript passed. Lint is clean of errors but has existing warnings outside this change; `format:check` exposed a local dependency gap: the `prettier` binary is missing from `node\_modules` despite being declared. I'm going to repair dependencies with `npm install` and rerun the check.

65. assistant (2026-07-09T08:44:11.083Z)

Confirmed: `prettier` is declared but absent locally. Installing dependencies should only hydrate `node\_modules`; I'll watch for lockfile churn and leave it out unless npm legitimately changes it.

66. assistant (2026-07-09T08:44:20.154Z)

Dependency hydration added the missing package without drama. I'm rerunning format and a quick status check before committing.

67. assistant (2026-07-09T08:44:28.544Z)

`format:check` now runs, but it fails across 30 existing `src` files, so that's repo-wide pre-existing formatting drift, not introduced by this PR. `npm install` did touch `package-lock.json`; I'm checking whether that lockfile change is meaningful before deciding whether to include or revert that single local artifact.

68. assistant (2026-07-09T08:44:40.937Z)

The targeted Gemini source file itself needs Prettier too. I'll format just that touched source file, then I'll revert the unrelated lockfile license churn with a tiny patch so the commit stays exactly on-topic.

69. assistant (2026-07-09T08:44:46.992Z)

The source file is formatted. I'm restoring the lockfile's pre-existing license value now so dependency installation doesn't sneak into the PR.

70. assistant (2026-07-09T08:44:55.074Z)

The lockfile is back to its original content, and the only intended files should now be the extractor and its regression tests. I'm doing one more verification sweep after formatting.

71. assistant (2026-07-09T08:45:00.503Z)

The touched Gemini source now passes Prettier. The lockfile is still marked modified, likely line endings rather than content, so I'm checking that exactly before I commit explicit paths.

72. assistant (2026-07-09T08:45:06.144Z)

There's no content diff in `package-lock.json`; Git is just seeing the working copy metadata/line-ending refresh. I'm refreshing the index rather than touching the file again.

73. assistant (2026-07-09T08:45:29.019Z)

The lockfile has no semantic diff, but npm left mixed line endings. I'm normalizing that one file back to the checkout's CRLF shape so the worktree is clean apart from the intended PR files.

74. assistant (2026-07-09T08:45:38.634Z)

The worktree is now scoped to the two intended files. I'm rerunning the full test command and lint once more after the cleanup, then I'll stage those exact paths and commit.

75. assistant (2026-07-09T08:45:48.021Z)

Verification is good: tests pass; lint has existing warnings only; the touched Gemini source passes Prettier. I'm doing the required branch/HEAD recheck, then I'll stage only the extractor and test file.

76. assistant (2026-07-09T08:45:52.090Z)

The branch is exactly on `origin/master` with no sibling commits on top. I'm staging the two intended files and committing the router P0.

77. assistant (2026-07-09T08:45:59.335Z)

Router commit is created. I'm checking the post-commit state, pushing the branch, and opening the second draft PR.

78. assistant (2026-07-09T08:46:02.026Z)

The branch contains only the intended commit and files. Pushing now.

79. assistant (2026-07-09T08:46:07.730Z)

Router branch is on GitHub. I'm opening it as a draft PR with the test and lint notes, including the repo-wide Prettier drift so there's no ambiguity.

80. assistant (2026-07-09T08:46:24.227Z)

Both draft PRs are open. Since this is a meaningful checkpoint, I'm doing the handoff refresh pass now so the next session sees the tracker and the two review PRs immediately.

81. assistant (2026-07-09T08:46:43.668Z)

I found a hub `SESSION\_HANDOFF.md` but no auto-memory handoff. I'm reading it before deciding whether a minimal refresh belongs in the already-open tracker PR or should stay local.

82. assistant (2026-07-09T08:46:59.004Z)

The hub handoff is still pointing at the July 7 baseline/replay work, so it's stale for this checkpoint. I'm replacing it with a concise current handoff that points to the feedback tracker and the two draft PRs, then I'll commit that documentation refresh into the hub PR.

83. assistant (2026-07-09T08:47:15.330Z)

Hub handoff is refreshed. I'm doing the branch/HEAD check again, then staging only `SESSION\_HANDOFF.md` and pushing that small follow-up to PR #68.

84. assistant (2026-07-09T08:47:18.128Z)

Hub branch still has only my PR commit on top of `origin/main`. I'm committing the handoff refresh as a separate small docs commit.

85. assistant (2026-07-09T08:47:24.958Z)

Hub PR #68 now includes the current handoff too. The router repo didn't have one, so I'm adding a short root `SESSION\_HANDOFF.md` there and pushing it to PR #18.

86. assistant (2026-07-09T08:47:40.322Z)

The router handoff is written. I'm checking branch/HEAD again, then I'll stage only that file and push the docs follow-up to the existing router PR.

87. assistant (2026-07-09T08:47:44.129Z)

Router branch still contains only my Gemini commit on top of `origin/master`. I'm committing the new handoff now.

88. assistant (2026-07-09T08:47:53.424Z)

Both PR branches are pushed with handoffs included. I'm doing one final status/PR check so I can give you clean links and the exact verification state.

89. assistant (2026-07-09T08:48:03.170Z)

Done. I opened both P0 draft PRs for your review:

- Hub PR: [avidullu/agent-sessions#68](https://github.com/avidullu/agent-sessions/pull/68)
Adds the feedback project tracker, refreshes the archive/PDF outputs, fixes archive index identity merging, ignores local router diagnostics, and updates the hub handoff.
- Router PR: [avidullu/agent-session-router#18](https://github.com/avidullu/agent-session-router/pull/18)
Sets stable Gemini Antigravity `metadata.session\_id` values from the `brain/<session>/system\_generated/logs/...` path, adds regression tests, and adds a router handoff.

Verification completed:

- Hub: `pytest` passed (`587 passed`, `96.81%` coverage), `ruff` passed, `mypy` passed.
- Router: `npm test` passed, `npm run lint:check` had existing warnings only, touched Gemini source passes Prettier.
- Both PR branches are clean and tracking their remotes.

One note: router-wide `npm run format:check` still fails across many existing `src` files, so I kept that out of the P0 PR scope and called it out in the PR body.

90. user (2026-07-09T08:48:03.563Z)

pls run the checked out agent-sessions indexer to discover and index and refresh/update teh archives as well

91. assistant (2026-07-09T08:48:10.912Z)

I'll refresh the hub archive from the currently checked-out `agent-sessions` branch, using the repo's own indexer/export path. First I'll re-sync safely with the remote branch, then run discover/index/export/status and commit any real archive changes back onto the PR branch.

92. assistant (2026-07-09T08:48:18.391Z)

The hub PR branch is up to date and clean. I'm checking the archive tool commands now so I use the project's intended discover/index/export flow instead of guessing at internal files.

93. assistant (2026-07-09T08:48:27.669Z)

The archive tool exposes ``discover``, ``export``, and ``status``; ``export`` is the operation that updates ``archive/index.jsonl`` and ``archive/INDEX.md``. I'm going to refresh discovery docs first, then run ``export --all --pdf`` so Markdown, PDFs, and indexes are all brought current.

94. assistant (2026-07-09T08:48:31.440Z)

The command set confirms the flow: ``discover --write``, then ``export --all --pdf``, then ``status --json``. I'm starting the actual refresh now; export may take a bit because it renders PDFs too.

95. assistant (2026-07-09T08:48:45.815Z)

Discovery refreshed ``docs/DISCOVERY.md``. I'm running the full archive export with PDF rendering now; this should also merge/update the archive indexes.